

serverless or capacity planning mentioned.

D. Configuring Amazon ElastiCache improves performance but does not directly address the operational overhead or capacity planning challenges mentioned.

Therefore, the correct answers are A and E as they address the requirements, while the incorrect answers (B, C, D) do not provide the desired solutions.

upvoted 3 times

studynoplay 2 years, 1 month ago

Selected Answer: AE

Improve the application's infrastructure = Modernize Infrastructure = Least Operational Overhead = Serverless

upvoted 4 times

Robrobtutu 2 years, 1 month ago

Selected Answer: AE

A and E are the best options.

upvoted 1 times

bgsanata 2 years, 3 months ago

Selected Answer: AE

A and E

upvoted 1 times

rapatajones 2 years, 4 months ago

Selected Answer: AE

a e.....

upvoted 1 times

goodmail 2 years, 5 months ago

One should that Aurora is not serverless. Aurora serverless and Aurora are 2 Amazon services. I prefer C, however the question does not mention any frontend requirements.

upvoted 1 times

aba2s 2 years, 5 months ago

Selected Answer: AE

Yes, go for A and E since the two resources are serverless.

upvoted 3 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: AE

The correct answers are A and E. To improve the application's infrastructure, the solutions architect should migrate the PostgreSQL database to Amazon Aurora and migrate the web application to be hosted on AWS Fargate with Amazon Elastic Container Service (Amazon ECS).

Amazon Aurora is a fully managed, scalable, and highly available relational database service that is compatible with PostgreSQL. Migrating the database to Amazon Aurora would reduce the operational overhead of maintaining the database infrastructure and allow the company to focus on building and scaling the application.

AWS Fargate is a fully managed container orchestration service that enables users to run containers without the need to manage the underlying EC2 instances. By using AWS Fargate with Amazon Elastic Container Service (Amazon ECS), the solutions architect can improve the scalability and efficiency of the web application and reduce the operational overhead of maintaining the underlying infrastructure.

upvoted 2 times

techhb 2 years, 5 months ago

A and E are obvious choices.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #130

Topic 1

An application runs on Amazon EC2 instances across multiple Availability Zones. The instances run in an Amazon EC2 Auto Scaling group behind an Application Load Balancer. The application performs best when the CPU utilization of the EC2 instances is at or near 40%.

What should a solutions architect do to maintain the desired performance across all instances in the group?

- A. Use a simple scaling policy to dynamically scale the Auto Scaling group.
- B. Use a target tracking policy to dynamically scale the Auto Scaling group. **Most Voted**
- C. Use an AWS Lambda function to update the desired Auto Scaling group capacity.
- D. Use scheduled scaling actions to scale up and scale down the Auto Scaling group.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: B

The correct answer is B. To maintain the desired performance across all instances in the Amazon EC2 Auto Scaling group, the solutions architect should use a target tracking policy to dynamically scale the Auto Scaling group.

A target tracking policy allows the Auto Scaling group to automatically adjust the number of EC2 instances in the group based on a target value for a metric. In this case, the target value for the CPU utilization metric could be set to 40% to maintain the desired performance of the application. The Auto Scaling group would then automatically scale the number of instances up or down as needed to maintain the target value for the metric.

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-scaling-simple-step.html>

upvoted 22 times

tom_cruise **Highly Voted** 1 year, 8 months ago

Selected Answer: B

target tracking policy = maintain

upvoted 6 times

Rcosmos **Most Recent** 5 months ago

Selected Answer: B

Justificativa:

O que é uma política de rastreamento de destino?

Uma política de rastreio de destino (Target Tracking Scaling Policy) ajusta automaticamente o tamanho do grupo de Auto Scaling para manter uma métrica-alvo específica, como a utilização da CPU, próxima a um valor desejado.

Nesse caso, a métrica-alvo seria 40% de utilização da CPU.

Por que a opção B é ideal?

Automação baseada em métricas:

O rastreamento de destino monitora continuamente a utilização da CPU das instâncias e ajusta a capacidade do grupo para manter o desempenho próximo ao alvo configurado (40%).

upvoted 1 times

satyaamm 5 months ago

Selected Answer: B

Since it is mentioned that it works best when CPU usage is around 40% hence using target tracking within ASG's is the most suitable option here.

upvoted 1 times

jaradat02 10 months, 3 weeks ago

Selected Answer: B

B is the correct answer, since target scaling monitors cloudwatch metrics, while simple/step scaling monitors cloudwatch alarms.

upvoted 3 times

1e22522 10 months, 2 weeks ago

Useful, jaradat02. Thanks.

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

40% CPU for best performance is a "target tracking" policy for scaling so B is correct.

A: Wrong policy

CD: Won't achieve 40% CPU

upvoted 3 times

Ruffyt 1 year, 7 months ago

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-scaling-simple-step.html>

upvoted 3 times

youdelin 1 year, 8 months ago

Selected Answer: B

I really don't get what kind of software running like a car with the most economical fuel speed range, but well, the answer is B

upvoted 2 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: B

The application performs best when the CPU utilization of the EC2 instances is at or near 40%.

Target tracking will maintain CPU utilization at 40%. When CloudWatch detects that the average CPU utilization is beyond 40%, it will trigger the target tracking policy to scale out the auto scaling group to meet this target utilization. Once everything is settled and the average CPU utilization has gone below 40%, another scale in action will kick in and reduce the number of auto scaling instances in the auto scaling group.

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

The key reasons are:

A target tracking policy allows defining a specific target metric value to maintain, in this case 40% CPU utilization.

Auto Scaling will automatically add or remove instances to keep utilization at the target level, without manual intervention.

This will dynamically scale the group to maintain performance as load changes.

This will dynamically scale the group to maintain performance as load changes.

A simple scaling policy only responds to breaching thresholds, not maintaining a target.

Scheduled actions and Lambda would require manual calculation and updates to track utilization.

Target tracking policies are the native Auto Scaling feature designed to maintain a metric at a target value.

upvoted 4 times

cookieMr 1 year, 11 months ago

Selected Answer: B

Target tracking policy is the most appropriate choice. This policy allows ASG to automatically adjust the desired capacity based on a target metric, such as CPU utilization. By setting the target metric to 40%, ASG will scale the number of instances up or down as needed to maintain the desired CPU utilization level. This ensures that the application's performance remains optimal.

A suggests using a simple scaling policy, which allows for scaling based on a fixed metric or threshold. However, it may not be as effective as a target tracking policy in dynamically adjusting the capacity to maintain a specific CPU utilization level.

C suggests using an Lambda to update the desired capacity. While this can be done programmatically, it would require custom scripting and may not provide the same level of automation and responsiveness as a target tracking policy.

D suggests using scheduled scaling actions to scale up and down ASG at predefined times. This approach is not suitable for maintaining the desired performance in real-time based on actual CPU utilization.

upvoted 3 times

Robrobtutu 2 years, 1 month ago

Selected Answer: B

B of course.

upvoted 1 times

aba2s 2 years, 5 months ago

Selected Answer: B

B seem to the correct response.

With a target tracking scaling policy, you can increase or decrease the current capacity of the group based on a target value for a specific metric. This policy will help resolve the over-provisioning of your resources. The scaling policy adds or removes capacity as required to keep the metric at, or close to, the specified target value. In addition to keeping the metric close to the target value, a target tracking scaling policy also adjusts to changes in the metric due to a changing load pattern.

upvoted 4 times

orionizzie 2 years, 5 months ago

Selected Answer: B

target tracking - CPU at 40%

upvoted 2 times

career360guru 2 years, 5 months ago

Selected Answer: B

Option B

upvoted 1 times

Wpcorgan 2 years, 6 months ago

B is correct

upvoted 1 times

ArielSchivo 2 years, 7 months ago

Selected Answer: B

Option B. Target tracking policy.

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-scaling-target-tracking.html>

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #131

Topic 1

A company is developing a file-sharing application that will use an Amazon S3 bucket for storage. The company wants to serve all the files through an Amazon CloudFront distribution. The company does not want the files to be accessible through direct navigation to the S3 URL.

What should a solutions architect do to meet these requirements?

- A. Write individual policies for each S3 bucket to grant read permission for only CloudFront access.
- B. Create an IAM user. Grant the user read permission to objects in the S3 bucket. Assign the user to CloudFront.
- C. Write an S3 bucket policy that assigns the CloudFront distribution ID as the Principal and assigns the target S3 bucket as the Amazon Resource Name (ARN).
- D. Create an origin access identity (OAI). Assign the OAI to the CloudFront distribution. Configure the S3 bucket permissions so that only the OAI has read permission. **Most Voted**

Correct Answer: D*Community vote distribution*

D (100%)

Comments**123jh10** **Highly Voted** 1 year, 7 months ago**Selected Answer: D**

I want to restrict access to my Amazon Simple Storage Service (Amazon S3) bucket so that objects can be accessed only through my Amazon CloudFront distribution. How can I do that?

Create a CloudFront origin access identity (OAI)

<https://aws.amazon.com/premiumsupport/knowledge-center/cloudfront-access-to-amazon-s3/>

upvoted 40 times

SimonPark 1 year, 7 months ago

Thanks it convinces me

upvoted 1 times

Guru4Cloud **Highly Voted** 10 months ago**Selected Answer: D**

The key reasons are:

An OAI provides secure access between CloudFront and S3 without exposing the S3 bucket publicly.

The OAI is associated with the CloudFront distribution.

The S3 bucket policy limits access only to that OAI.

This ensures only CloudFront can access the objects, not direct S3 access.

Option A is complex to manage individual bucket policies.

Option B exposes credentials that aren't needed.

Option C works but OAI is the preferred method.

So using an origin access identity provides the most secure way to serve private S3 content through CloudFront. The OAI prevents direct public access to the S3 bucket.

upvoted 10 times

satyaamm Most Recent 5 months ago

Selected Answer: D

OAI's are designed specifically for S3 origins. It helps provide security for CloudFront.

upvoted 1 times

xdkonorek2 7 months, 1 week ago

Selected Answer: D

C would also work but missing important details in the answer

D is legacy and architect should not recommend it

upvoted 4 times

tom_cruise 8 months ago

Selected Answer: D

"If your users try to access objects using Amazon S3 URLs, they're denied access. The origin access identity has permission to access objects in your Amazon S3 bucket, but users don't."

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/private-content-restricting-access-to-s3.html>

upvoted 2 times

cookieMr 11 months, 3 weeks ago

Selected Answer: D

To meet the requirements of serving files through CloudFront while restricting direct access to the S3 bucket URL, the recommended approach is to use an origin access identity (OAI). By creating an OAI and assigning it to the CloudFront distribution, you can control access to the S3 bucket.

This setup ensures that the files stored in the S3 bucket are only accessible through CloudFront and not directly through the S3 bucket URL. Requests made directly to the S3 URL will be blocked.

Option A suggests writing individual policies for each S3 bucket, which can be cumbersome and difficult to manage, especially if there are multiple buckets involved.

Option B suggests creating an IAM user and assigning it to CloudFront, but this does not address restricting direct access to the S3 bucket URL.

Option C suggests writing an S3 bucket policy with CloudFront distribution ID as the Principal, but this alone does not provide the necessary restrictions to prevent direct access to the S3 bucket URL.

upvoted 4 times

antropaws 1 year ago

DECEMBER 2022 UPDATE:

Restricting access to an Amazon S3 origin:

CloudFront provides two ways to send authenticated requests to an Amazon S3 origin: origin access control (OAC) and origin access identity (OAI). We recommend using OAC because it supports:

All Amazon S3 buckets in all AWS Regions, including opt-in Regions launched after December 2022

Amazon S3 server-side encryption with AWS KMS (SSE-KMS)

Dynamic requests (PUT and DELETE) to Amazon S3

OAI doesn't work for the scenarios in the preceding list, or it requires extra workarounds in those scenarios.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/private-content-restricting-access-to-s3.html>

upvoted 2 times

Buruguduystunstugudunstuy 1 year, 5 months ago

Selected Answer: D

The correct answer is D. To meet the requirements, the solutions architect should create an origin access identity (OAI) and assign it to the CloudFront distribution. The S3 bucket permissions should be configured so that only the OAI has read permission.

An OAI is a special CloudFront user that is associated with a CloudFront distribution and is used to give CloudFront access to the files in an S3 bucket. By using an OAI, the company can serve the files through the CloudFront distribution while preventing direct access to the S3 bucket.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/private-content-restricting-access-to-s3.html>

upvoted 6 times

career360guru 1 year, 5 months ago

Selected Answer: D

D is the right answer

upvoted 1 times

gloritown 1 year, 6 months ago

Selected Answer: D

D is correct but instead of OAI using OAC would be better since OAI is legacy

upvoted 3 times

Robrobtutu 1 year, 1 month ago

Thanks, I didn't know about OAC.

upvoted 1 times

Wpcorgan 1 year, 6 months ago

D is correct

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #132

Topic 1


A company's website provides users with downloadable historical performance reports. The website needs a solution that will scale to meet the company's website demands globally. The solution should be cost-effective, limit the provisioning of infrastructure resources, and provide the fastest possible response time.

Which combination should a solutions architect recommend to meet these requirements?

- A. Amazon CloudFront and Amazon S3 **Most Voted**
- B. AWS Lambda and Amazon DynamoDB
- C. Application Load Balancer with Amazon EC2 Auto Scaling
- D. Amazon Route 53 with internal Application Load Balancers

Correct Answer: A

Community vote distribution



Comments

G3 **Highly Voted** 2 years, 4 months ago

Selected Answer: A

Historical reports = Static content = S3
upvoted 21 times

dokaedu **Highly Voted** 2 years, 7 months ago

A is the correct answer
The solution should be cost-effective, limit the provisioning of infrastructure resources, and provide the fastest possible response time.
upvoted 11 times

Manimgh **Most Recent** 2 months ago

Selected Answer: A

Question mentioned cost effective and S3 is the most cost effective answer
upvoted 1 times

PaulGa 8 months, 4 weeks ago

Selected Answer: A

Ans A - S3 is designed to optimise costs, highly scalable and can store static content such as website. CloudFront is designed to securely deliver content with low latency and high transfer rate.

upvoted 2 times

huaze_lei 9 months, 1 week ago

Selected Answer: A

Historical data will not change, and hence they are static content. So the answer is S3 with distributed content (Cloudfront)

upvoted 2 times

jaradat02 10 months, 3 weeks ago

Selected Answer: A

A is the most suitable choice because the content is static and downloadable

upvoted 2 times

KTEgghead 11 months, 1 week ago

Selected Answer: A

A. S3 is designed to optimize storage costs, is highly scalable and can hold static content e.g. a website. CloudFront is designed to securely deliver content with low latency and high transfer speeds.

upvoted 2 times

MrPCarrot 1 year, 5 months ago

Bringing content closer to users, Answer is A

upvoted 3 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: A

Global, cost-effective, serverless, low latency = CloudFront with S3

Static content = S3

upvoted 6 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: A

Historical reports = Static content = S3

upvoted 3 times

cookieMr 1 year, 11 months ago

By using CloudFront, the website can leverage the global network of edge locations to cache and deliver the performance reports to users from the nearest edge location, reducing latency and providing fast response times. Amazon S3 serves as the origin for the files, where the reports are stored.

Option B is incorrect because AWS Lambda and Amazon DynamoDB are not the most suitable services for serving downloadable files and meeting the website demands globally.

Option C is incorrect because using an Application Load Balancer with Amazon EC2 Auto Scaling may require more infrastructure provisioning and management compared to the CloudFront and S3 combination. Additionally, it may not provide the same level of global scalability and fast response times as CloudFront.

Option D is incorrect because while Amazon Route 53 is a global DNS service, it alone does not provide the caching and content delivery capabilities required for serving the downloadable reports. Internal Application Load Balancers do not address the global scalability and caching requirements specified in the scenario.

upvoted 5 times

Bmarodi 1 year, 11 months ago

Very good explanations!

upvoted 1 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: A

This is a question about AWS services. The question is asking for the best solution to deliver static content globally with low latency and high transfer rates. The options are: A. Amazon S3 and Amazon CloudFront, B. Amazon S3 and Amazon ElastiCache, C. Amazon S3 and Amazon EC2, D. Amazon S3 and Amazon Route 53. The correct answer is A. Amazon S3 and Amazon CloudFront. Amazon S3 is a scalable object storage service, and Amazon CloudFront is a content delivery network (CDN) service. Together, they provide a highly scalable and secure way to deliver static content globally with low latency and high transfer rates. Amazon ElastiCache is a managed in-memory cache service, Amazon EC2 is a virtual machine service, and Amazon Route 53 is a DNS service. None of these services are designed for delivering static content globally with low latency and high transfer rates.

The correct answer is Option A. To meet the requirements, the solutions architect should recommend using Amazon CloudFront and Amazon S3.

By combining Amazon CloudFront and Amazon S3, the solutions architect can provide a scalable and cost-effective solution that limits the provisioning of infrastructure resources and provides the fastest possible response time.

<https://aws.amazon.com/cloudfront/>

<https://aws.amazon.com/s3/>

upvoted 4 times

techhb 2 years, 5 months ago

A is correct

upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: A

A is the best and most cost effective option if only download of the static pre-created report(no data processing before downloading) is a requirement.

upvoted 2 times

Wpcorgan 2 years, 6 months ago

A is correct

upvoted 1 times

sdasdawa 2 years, 7 months ago

Selected Answer: A

<https://www.examttopics.com/discussions/amazon/view/27935-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Nirmal3331 2 years, 7 months ago

Selected Answer: A

<https://www.examttopics.com/discussions/amazon/view/27935-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #133

Topic 1

A company runs an Oracle database on premises. As part of the company's migration to AWS, the company wants to upgrade the database to the most recent available version. The company also wants to set up disaster recovery (DR) for the database. The company needs to minimize the operational overhead for normal operations and DR setup. The company also needs to maintain access to the database's underlying operating system.

Which solution will meet these requirements?

- A. Migrate the Oracle database to an Amazon EC2 instance. Set up database replication to a different AWS Region.
- B. Migrate the Oracle database to Amazon RDS for Oracle. Activate Cross-Region automated backups to replicate the snapshots to another AWS Region.
- C. Migrate the Oracle database to Amazon RDS Custom for Oracle. Create a read replica for the database in another AWS Region. **Most Voted**
- D. Migrate the Oracle database to Amazon RDS for Oracle. Create a standby database in another Availability Zone.

Correct Answer: C*Community vote distribution***Comments****ArielSchivo** **Highly Voted** 2 years, 7 months ago

Option C since RDS Custom has access to the underlying OS and it provides less operational overhead. Also, a read replica in another Region can be used for DR activities.

<https://aws.amazon.com/blogs/database/implementing-a-disaster-recovery-strategy-with-amazon-rds/>
upvoted 45 times

KalarAzar 1 year, 12 months ago

You can't create cross-Region replicas in RDS Custom for Oracle:
<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/custom-rr.html#custom-rr.limitations>
upvoted 23 times

brushek **Highly Voted** 2 years, 8 months ago**Selected Answer: C**

It should be C:

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/rds-custom.html>

and

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/working-with-custom-oracle.html>

upvoted 21 times

bhgt 1 year, 8 months ago

how it is C when the read replica is not meant for DR

upvoted 3 times

wsdasdasdqwda 1 year, 7 months ago

If the source DB instance fails, you can promote your Read Replica to a standalone source server.

upvoted 11 times

CloudExpert01 Most Recent 2 months, 1 week ago

Selected Answer: C

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/working-with-custom-oracle.html>

upvoted 1 times

AshishDhole 4 months ago

Selected Answer: C

The company also needs to maintain access to the database's underlying operating system --> only RDS Custom can provide this facility

upvoted 1 times

zdi561 4 months, 2 weeks ago

Selected Answer: C

you can create a read replica in another region which is less operational cost for DR

upvoted 1 times

Ode7d1b 6 months, 3 weeks ago

Selected Answer: C

Amazon RDS Custom for Oracle:

Amazon RDS Custom for Oracle provides a managed database service while also allowing access to the underlying operating system, which is required in this scenario.

This option strikes a balance between minimizing operational overhead (as AWS manages backups, patching, and scaling) and retaining flexibility for customizations or specific access requirements.

upvoted 2 times

TheTeaBoy 7 months ago

A: Incorrect - Question states to minimise operational overhead, building an EC2 instance and managing that doesn't reduce the operational overhead.

B: Incorrect - Yes that ticks the DR aspect, but Amazon RDS for Oracle doesn't provide access to the underlying operating system of the database server.

C: Correct - Amazon RDS Custom provides a limited amount of operating system access of the database server(s), so that's good. Creating a read replica in another AWS Region covers DR. Its not asking for High Availability (HA) in this case.

D: Incorrect - Amazon RDS for Oracle doesn't provide access to the underlying operating system of the database server. A standby database in another availability, could be considered as covering DR.

upvoted 3 times

mzeynalli 7 months, 3 weeks ago

Selected Answer: C

Amazon RDS Custom for Oracle is the best choice because it allows access to the underlying OS, provides an easier path for upgrades, minimizes operational overhead, and supports setting up cross-region read replicas for a DR solution.

upvoted 2 times

19d92c7 7 months, 3 weeks ago

It should be standby

it should be standby.

. When the primary database fails,

Amazon RDS promotes the secondary database to primary. Because it assumes the primary databases endpoint, the EC2 instances can resume traffic with the new primary database. Meanwhile, a new standby database is created in the other Availability Zone.

upvoted 1 times

PaulGa 8 months, 4 weeks ago

Selected Answer: D

Ans D - use regular RDS Oracle; no need for RDS Custom because the question doesn't state and special customisations... the only bit unanswered is "...maintain access to company's underlying o/s"

upvoted 2 times

LuongTo 7 months, 3 weeks ago

D is out since DR is more about "different region", not just "different AZ"

upvoted 2 times

bignatov 9 months, 2 weeks ago

Selected Answer: C

Option C provides the right balance of managed service convenience, access to the underlying OS, and effective disaster recovery with cross-region read replicas

upvoted 2 times

MatAlves 9 months, 2 weeks ago

Selected Answer: A

"You can't create cross-Region RDS Custom for Oracle replicas." So it can't be C.

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/custom-rr.html>

B and D don't provide access to underlying OS. So the only option left is A, which won't help minimize the operational overhead though.

upvoted 4 times

scaredSquirrel 10 months ago

Selected Answer: C

B, D - do not have access to underlying OS.

Both A and C could work, but C is less Overhead.

C - Amazon RDS Custom for Oracle actually supports creating read replicas. "Creating an RDS Custom for Oracle replica is similar to creating an RDS for Oracle replica, but with important differences."

(<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/custom-rr.html>). The read replica is not meant for disaster recovery, but it could work as such when no better options are available.

upvoted 2 times

MatAlves 9 months, 2 weeks ago

"You can't create cross-Region RDS Custom for Oracle replicas."

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/custom-rr.html>

upvoted 1 times

Abbas_Abi_AWS 10 months, 1 week ago

Selected Answer: A

You can't create cross-Region RDS Custom for Oracle replicas.

upvoted 3 times

jaradat02 10 months, 3 weeks ago

Selected Answer: C

Since we need to have access to the underlying infrastructure, C makes sense.

upvoted 1 times

ChinthaGurumurthi 10 months, 3 weeks ago

Selected Answer: C

Selected Answer: C

Its C

Key : Access to underlying O.S - RDS custom can give you this feature

General RDS for Oracle(or any) - you can't access underlying O.S

So, definitely C considering this point in the question

upvoted 1 times

freedafeng 11 months, 1 week ago

Definitely not C. You cannot create a read replica of RDS custom for Oracle in a different region:

<https://aws.amazon.com/blogs/database/build-high-availability-for-amazon-rds-custom-for-oracle-using-read-replicas/>

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #134

Topic 1

A company wants to move its application to a serverless solution. The serverless solution needs to analyze existing and new data by using SL. The company stores the data in an Amazon S3 bucket. The data requires encryption and must be replicated to a different AWS Region.


Which solution will meet these requirements with the LEAST operational overhead?

A. Create a new S3 bucket. Load the data into the new S3 bucket. Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region. Use server-side encryption with AWS KMS multi-Region keys (SSE-KMS). Use Amazon Athena to query the data.

B. Create a new S3 bucket. Load the data into the new S3 bucket. Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region. Use server-side encryption with AWS KMS multi-Region keys (SSE-KMS). Use Amazon RDS to query the data.

C. Load the data into the existing S3 bucket. Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region. Use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Use Amazon Athena to query the data. **Most Voted**

D. Load the data into the existing S3 bucket. Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region. Use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Use Amazon RDS to query the data.

Correct Answer: C*Community vote distribution***Comments****123jh10** **Highly Voted** 2 years, 7 months ago**Selected Answer: C**

SSE-KMS vs SSE-S3 - The last seems to have less overhead (as the keys are automatically generated by S3 and applied on data at upload, and don't require further actions. KMS provides more flexibility, but in turn involves a different service, which finally is more "complex" than just managing one (S3). So A and B are excluded. If you are in doubt, you are having 2 buckets in A and B, while just keeping one in C and D.

<https://s3browser.com/server-side-encryption-types.aspx>

Decide between C and D is deciding on Athena or RDS. RDS is a relational db, and we have documents on S3, which is the use case for Athena. Athena is also serverless, which eliminates the need of controlling the underlying infrastructure and capacity. So C is the answer.

so C is the answer.

<https://aws.amazon.com/athena/>

upvoted 65 times

markw92 1 year, 12 months ago

See comment from Nicknameinvalid below. You get your answer.

upvoted 2 times

MutiverseAgent 1 year, 11 months ago

It's since replication works for new objects but not for the existing ones, unless you use batch replication which is not the case.

upvoted 1 times

Chiznitz 1 year, 7 months ago

Answer A has you move the data before you enable replication, therefore there is no difference between A and C when it comes to the point in time you enable replication. I agree A would be a better choice if the order of operations said, create a bucket->Enable encryption->move files...but it doesn't. It has you create the bucket and move the files.

upvoted 3 times

dokaedu Highly Voted 2 years, 7 months ago

Answer is A:

Amazon S3 Bucket Keys reduce the cost of Amazon S3 server-side encryption using AWS Key Management Service (SSE-KMS). This new bucket-level key for SSE can reduce AWS KMS request costs by up to 99 percent by decreasing the request traffic from Amazon S3 to AWS KMS. With a few clicks in the AWS Management Console, and without any changes to your client applications, you can configure your bucket to use an S3 Bucket Key for AWS KMS-based encryption on new objects. The Existing S3 bucket might have unencrypted data - encryption will apply new data received after the applying of encryption on the new bucket.

upvoted 30 times

AKBM7829 1 year, 9 months ago

But in server side encryption Multi Region Keys is not possible which leaves to Option C

upvoted 3 times

NSA_Poker 1 year ago

"you manage the multi-Region key in each Region independently. Neither AWS nor AWS KMS ever automatically creates or replicates multi-Region keys into any Region on your behalf. AWS managed keys, the KMS keys that AWS services create in your account for you, are always single-Region keys."

<https://docs.aws.amazon.com/kms/latest/developerguide/multi-region-keys-overview.html>

upvoted 1 times

SAgang 11 months, 1 week ago

from your link you missed this part:

you can encrypt data in one AWS Region and decrypt it in a different AWS Region without re-encrypting or making a cross-Region call to AWS KMS

upvoted 1 times

s50600822 2 years, 1 month ago

Don't know what "kays" are, could they be a trap?

upvoted 1 times

Bmarodi 2 years ago

Kays = keys, mistype i think.

upvoted 1 times

RBSK 2 years, 6 months ago

Cost reduction is in comparison bet Bucket level KMS key and object level KMS key. Not between SSE-KMS and SSE-S3. Hence its a wrong comparison

upvoted 2 times

RODROSKAR 2 years, 7 months ago

Reducing cost was never the target, it's LEAST operational. In that regard SSE-S3 AWS fully managed.

upvoted 6 times

MortisG Most Recent 3 weeks, 1 day ago

Selected Answer: A

Correct answer is A. SSE-S3 is not supported for CRR.

upvoted 1 times

CloudExpert01 2 months, 1 week ago

Selected Answer: C

since the requirement is "The serverless solution needs to analyze existing and new data", the answer C seems to be more appropriate since all the data will be in 1 bucket

upvoted 1 times

Mimine87 2 months, 1 week ago

Selected Answer: A

Amazon S3 + Athena is a fully serverless data analytics solution (no infrastructure to manage).

S3 Cross-Region Replication (CRR) enables automatic, asynchronous replication of objects across AWS Regions, satisfying the replication requirement.

SSE-KMS with multi-Region keys ensures that the data is encrypted and can be decrypted in the destination Region, which is crucial for cross-region analytics and compliance.

Athena is used for analyzing data directly in S3 using SQL, which meets the requirement to analyze existing and new data using SQL (SL was a typo).

This option delivers everything with the least operational overhead — no server provisioning, no DB management, and built-in encryption + replication.

upvoted 1 times

Faraz999 2 months, 2 weeks ago

Selected Answer: A

you can apply cross-region replication to an existing Amazon S3 bucket, but it will only replicate new objects after the replication rule is configured. To replicate existing objects, you'll need to use S3 Batch Replication.

upvoted 1 times

MPG1970 2 months, 2 weeks ago

Selected Answer: A

I went for A because, if you start to encrypt only new data then

1. You will have an inconsistently solution for security. I think it fair to assume that if they want all new data encrypted, then all data should be encrypted
2. the serverless function will become more complex as it will have to establish whether the data is encrypted or not every time it retrieves an object.

Therefore you need to create a new S3 bucket so that all data will be encrypted

upvoted 1 times

SirDNS 2 months, 3 weeks ago

Selected Answer: C

fewer services = less operational overhead

managed services/serverless = less operational overhead

upvoted 1 times

tch 3 months ago

Selected Answer: A

If you want a straightforward approach to encrypting your S3 data without managing your own encryption keys, SSE-S3 is a good option.

you need AWS KMS for this complex solution
upvoted 1 times

zdi561 4 months, 1 week ago

Selected Answer: A

A new bucket is needed to encrypt objects. C uses existing bucket.
upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: A

About typos: "kays" should be "keys", and "SL" should be "SQL". So you gotta choose between A and C (We need Athena!).

About option C: SSE-S3 is a valid encryption method, just not as suitable as SSE-KMS with multi-Region keys for CRR. SSE-KMS with multi-Region keys simplifies key management in the destination Region.
upvoted 1 times

Rcosmos 5 months ago

Selected Answer: U

A melhor opção para atender aos requisitos com a menor sobrecarga operacional é a opção A:

A. Crie um novo bucket do S3. Carregue os dados no novo bucket do S3. Use a replicação entre regiões (CRR) do S3 para replicar objetos criptografados para um bucket do S3 em outra região. Use a criptografia no lado do servidor com chaves multirregionais do AWS KMS (SSE-KMS). Use o Amazon Athena para consultar os dados. Conclusão:
A opção A é a escolha ideal, pois combina os recursos sem servidor do Amazon Athena, a segurança avançada do SSE-KMS com chaves multirregionais, e a replicação automática entre regiões com o S3 CRR, tudo com a menor sobrecarga operacional.
upvoted 1 times

skyerwhite 5 months, 2 weeks ago

Selected Answer: A

what is the different between create a new S3 bucket and load data into the existing S3 bucket? I don't get it this. Please..
upvoted 1 times

Mischi 5 months, 3 weeks ago

Selected Answer: A

Option A provides a serverless solution, advanced encryption with AWS KMS multi-region keys and cross-region replication with CRR, all with the lowest operational overhead. Amazon Athena is the ideal tool for analyzing data in S3 without the need for additional infrastructure.
upvoted 1 times

Chr1s_Mrg 8 months, 1 week ago

Selected Answer: C

RDS is relational DB so we need Athena for this
upvoted 1 times

tonybuivannghia 8 months, 3 weeks ago

Selected Answer: A

I think A is correct because SSE-S3 doesn't support multi-region key management, but SSE-KMS has.
upvoted 1 times

PaulGa 8 months, 4 weeks ago

Selected Answer: A

Ans A - once you realise "SL" is a typo for "ML" then its only the Athena options, and in the case of option it means setting up a new S3 bucket
upvoted 3 times

Ilccing 5 months, 2 weeks ago

I'm curious that if the SL is the typo of "SQL"?
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #135

Topic 1

A company runs workloads on AWS. The company needs to connect to a service from an external provider. The service is hosted in the provider's VPC. According to the company's security team, the connectivity must be private and must be restricted to the target service. The connection must be initiated only from the company's VPC.

Which solution will meet these requirements?

- A. Create a VPC peering connection between the company's VPC and the provider's VPC. Update the route table to connect to the target service.
- B. Ask the provider to create a virtual private gateway in its VPC. Use AWS PrivateLink to connect to the target service.
- C. Create a NAT gateway in a public subnet of the company's VPC. Update the route table to connect to the target service.
- D. Ask the provider to create a VPC endpoint for the target service. Use AWS PrivateLink to connect to the target service.

Most Voted**Correct Answer: D***Community vote distribution*

D (100%)

Comments**123jhl0** **Highly Voted** 2 years, 7 months ago**Selected Answer: D**

****AWS PrivateLink provides private connectivity between VPCs, AWS services, and your on-premises networks, without exposing your traffic to the public internet**.** AWS PrivateLink makes it easy to connect services across different accounts and VPCs to significantly simplify your network architecture.

Interface ****VPC endpoints****, powered by AWS PrivateLink, connect you to services hosted by AWS Partners and supported solutions available in AWS Marketplace.

<https://aws.amazon.com/privatelink/>

upvoted 33 times

remand **Highly Voted** 2 years, 4 months ago**Selected Answer: D**

The solution that meets these requirements best is option D.

By asking the provider to create a VPC endpoint for the target service, the company can use AWS PrivateLink to connect to the target service. This enables the company to access the service privately and securely over an Amazon VPC endpoint, without requiring a NAT gateway, VPN, or AWS Direct Connect. Additionally, this will restrict the connectivity only to the target service.

requiring a NAT gateway, VPN, or AWS Direct Connect. Additionally, this will restrict the connectivity only to the target service, as required by the company's security team.

Option A VPC peering connection may not meet security requirement as it can allow communication between all resources in both VPCs.

Option B, asking the provider to create a virtual private gateway in its VPC and use AWS PrivateLink to connect to the target service is not the optimal solution because it may require the provider to make changes and also you may face security issues. Option C, creating a NAT gateway in a public subnet of the company's VPC can expose the target service to the internet, which would not meet the security requirements.

upvoted 11 times

PaulGa Most Recent 8 months, 4 weeks ago

Selected Answer: D

Ans D - create a unique, private only link: "Ask the provider to create a VPC endpoint for the target service. Use AWS PrivateLink to connect to the target service"

upvoted 2 times

lofzee 1 year ago

Selected Answer: D

no split decisions on this answer eh? not like the last one. lol

upvoted 2 times

RNess 1 year, 7 months ago

Selected Answer: D

AWS PrivateLink / VPC Endpoint Services:

- Connect services privately from your service VPC to customers VPC
- Doesn't need VPC Peering, public Internet, NAT Gateway, Route Tables
- Must be used with Network Load Balancer & ENI

upvoted 3 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: D

option D is correct

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: D

The best solution to meet the requirements is option D:

Ask the provider to create a VPC endpoint for the target service

Use AWS PrivateLink to connect to the target service

The reasons are:

PrivateLink provides private connectivity between VPCs without using public internet.

The provider creates a VPC endpoint in their VPC for the target service.

The company uses PrivateLink to securely access the endpoint from their VPC.

Connectivity is restricted only to the target service.

The connection is initiated only from the company's VPC.

Options A, B, C would expose the connection to the public internet or require infrastructure changes in the provider's VPC.

PrivateLink enables private, restricted connectivity to the target service without VPC peering or public exposure.

upvoted 3 times

cookieMr 1 year, 11 months ago

Selected Answer: D

Option C meets the requirements of establishing a private and restricted connection to the service hosted in the provider's VPC. By asking the provider to create a VPC endpoint for the target service, you can establish a direct and private connection from your company's VPC to the target service. AWS PrivateLink ensures that the connectivity remains within the AWS network and does not require internet access. This ensures both privacy and restriction to the target service, as the connection can only be initiated from your company's VPC.

A. VPC peering does not restrict access only to the target service.

B. PrivateLink is typically used for accessing AWS services, not external services in a provider's VPC.

C. NAT gateway does not provide a private and restricted connection to the target service.

Option D is the correct choice as it uses AWS PrivateLink and VPC endpoint to establish a private and restricted connection from the company's VPC to the target service in the provider's VPC.

upvoted 4 times

Abrar2022 2 years ago

VPC Endpoint (Target Service) - for specific services (not accessing whole vpc)

VPC Peering - (accessing whole VPC)

upvoted 4 times

Abrar2022 2 years ago

VPC Peering Connection:

All resources in a VPC, such as ECSs and load balancers, can be accessed.

VPC Endpoint:

Allows access to a specific service or application. Only the ECSs and load balancers in the VPC for which VPC endpoint services are created can be accessed.

upvoted 2 times

eugene_stalker 2 years ago

Selected Answer: D

Option D, but seems that it is vice versa. Customer needs to create PrivateLink and and you VPC endpoint to connect to PrivateLink

upvoted 2 times

studynoplay 2 years, 1 month ago

AWS PrivateLink / VPC Endpoint Services:

- Connect services privately from your service VPC to customers VPC
- Doesn't need VPC Peering, public Internet, NAT Gateway, Route Tables
- Must be used with Network Load Balancer & ENI

upvoted 3 times

Help2023 2 years, 3 months ago

Selected Answer: D

D. Here you are the one initiating the connection

upvoted 2 times

devonwho 2 years, 4 months ago

Selected Answer: D

PrivateLink is a more generalized technology for linking VPCs to other services. This can include multiple potential endpoints: AWS services, such as Lambda or EC2; Services hosted in other VPCs; Application endpoints hosted on-premises.

<https://www.tinystacks.com/blog-post/aws-vpc-peering-vs-privatelink-which-to-use-and-when/>

upvoted 2 times

devonwho 2 years, 4 months ago

Selected Answer: D

While VPC peering enables you to privately connect VPCs, AWS PrivateLink enables you to configure applications or services in VPCs as endpoints that your VPC peering connections can connect to.

upvoted 3 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: D

The solution that meets these requirements is Option D:

- * Ask the provider to create a VPC endpoint for the target service.
- * Use AWS PrivateLink to connect to the target service.

Option D involves asking the provider to create a VPC endpoint for the target service, which is a private connection to the service that is hosted in the provider's VPC. This ensures that the connection is private and restricted to the target service, as required by the company's security team. The company can then use AWS PrivateLink to connect to the target service over the VPC endpoint. AWS PrivateLink is a fully managed service that enables you to privately access services hosted on AWS or on-premises.

VPC endpoint. AWS PrivateLink is a fully managed service that enables you to privately access services hosted on AWS, on-premises, or in other VPCs. It provides secure and private connectivity to services by using private IP addresses, which ensures that traffic stays within the Amazon network and does not traverse the public internet.

Therefore, Option D is the solution that meets the requirements.

upvoted 5 times

Buruguduystunstugudunstuy 2 years, 5 months ago

AWS PrivateLink documentation: <https://docs.aws.amazon.com/privatelink/latest/userguide/what-is-privatelink.html>

upvoted 2 times

techhb 2 years, 5 months ago

D is right, if requirement was to be ok with public internet then option C was ok.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #136

Topic 1


A company is migrating its on-premises PostgreSQL database to Amazon Aurora PostgreSQL. The on-premises database must remain online and accessible during the migration. The Aurora database must remain synchronized with the on-premises database.

Which combination of actions must a solutions architect take to meet these requirements? (Choose two.)

- A. Create an ongoing replication task. **Most Voted**
- B. Create a database backup of the on-premises database.
- C. Create an AWS Database Migration Service (AWS DMS) replication server. **Most Voted**
- D. Convert the database schema by using the AWS Schema Conversion Tool (AWS SCT).
- E. Create an Amazon EventBridge (Amazon CloudWatch Events) rule to monitor the database synchronization.

Correct Answer: AC

Community vote distribution



Comments

123jh10 **Highly Voted** 2 years, 7 months ago

Selected Answer: AC

AWS Database Migration Service (AWS DMS) helps you migrate databases to AWS quickly and securely. The source database remains fully operational during the migration, minimizing downtime to applications that rely on the database.

... With AWS Database Migration Service, you can also continuously replicate data with low latency from any supported source to any supported target.

<https://aws.amazon.com/dms/>

upvoted 33 times

pentium75 1 year, 5 months ago

PostgreSQL -> Aurora PostgreSQL requires schema conversion per <https://aws.amazon.com/dms/schema-conversion-tool/>

upvoted 3 times

LoXoL 1 year, 5 months ago

SCT is compatible with PostgreSQL as source and Aurora PostgreSQL as destination, but not required.

upvoted 3 times

gustavtd Highly Voted 2 years, 5 months ago

Selected Answer: AC

AC, here it is clearly shown https://docs.aws.amazon.com/zh_cn/dms/latest/sbs/chap-manageddatabases.postgresql-rds-postgresql.html
upvoted 11 times

LuckyAro 2 years, 4 months ago

You nailed it !
upvoted 1 times

Ode7d1b Most Recent 6 months, 3 weeks ago

Selected Answer: AC

DMS helps to migrate the data from onpremise to aws and require replication task
upvoted 1 times

PaulGa 8 months, 4 weeks ago

Selected Answer: CE

Ans C, E -
C: AWS Database Migration Service to migrate databases to AWS, source database remains fully operational during the migration, avoiding application downtime.
<https://aws.amazon.com/dms/>
E: monitor with CloudWatch
As for C - not convinced: its PostgreSQL to PostgreSQL migration... no SCT is needed?
upvoted 1 times

MatAlves 8 months, 3 weeks ago

No, since you're using the same schema for Postgresql.
upvoted 3 times

jaradat02 10 months, 3 weeks ago

Selected Answer: AC

A and c obviously
upvoted 1 times

jatric 11 months, 1 week ago

Selected Answer: AC

AC is more accurate as DMS will help to migrate the on-premises data base to cloud with ease and for ongoing replication to synchronized the database "ongiong replication task" will be helpfull.

And yes its PostgreSQL to PostgreSQL migration so no SCT is needed ehre
upvoted 2 times

BombArat 11 months, 3 weeks ago

Selected Answer: AC

WeLL CHATGPT says SCT is not required so , AC makes sense
upvoted 2 times

CCCat 1 year ago

Selected Answer: AC

Keywords:

- migrating its on-premises PostgreSQL database to Amazon Aurora PostgreSQL
- The Aurora database must remain synchronized with the on-premises database

Analysis:

Option A satisfy the requirement of "synchronized with the on-premises database"

Option C suits for the homogeneous database migration.

Option D is not needed in this scenario, it suits for the heterogeneous database.

Homogeneous database migration tools: <https://docs.aws.amazon.com/prescriptive-guidance/latest/migration-oracle->

homogeneous database migration tools: <https://docs.aws.amazon.com/prescriptive-guidance/latest/migration-tools/database/homogeneous-migration-tools.html>

Heterogeneous database migration tools: <https://aws.amazon.com/dms/schema-conversion-tool/>

upvoted 2 times

[Removed] 1 year ago

CD:

Migrating a schema from PostgreSQL to Amazon Aurora (PostgreSQL) usually requires using the AWS Schema Conversion Tool (SCT) and the AWS Database Migration Service (DMS)

upvoted 1 times

Hopeyemi 1 year, 1 month ago

To meet the requirements of migrating an on-premises PostgreSQL database to Amazon Aurora PostgreSQL while keeping the on-premises database online and ensuring synchronization with the Aurora database, the following actions need to be taken:

Create an ongoing replication task (Option A): This action involves setting up continuous replication between the on-premises PostgreSQL database and the Aurora PostgreSQL database. This ensures that changes made to the on-premises database are replicated to the Aurora database in real-time, keeping them synchronized.

Create an AWS Database Migration Service (AWS DMS) replication server (Option C): AWS DMS provides a reliable and efficient way to migrate databases to AWS while minimizing downtime. By creating an AWS DMS replication server, you can configure and manage the replication tasks between the on-premises database and the Aurora database.

upvoted 4 times

Alphateccc 1 year, 3 months ago

answer is CD : postgresql and aurora postgresql have different schemes, you need sct for conversion and dms for the migration (replication)

upvoted 1 times

vip2 1 year, 3 months ago

Selected Answer: AC

AC

perform ongoing replication using AWS DMS to keep the source and target databases in sync

upvoted 2 times

farnamjam 1 year, 5 months ago

Selected Answer: AC

DMS has Continuous Data Replication using CDC

upvoted 3 times

Michael_Li 1 year, 6 months ago

CD

A is out because it does not specify what is the service to perform the replication task, clearly what needed here is DMS

B is out because backup is solution to keep 2 DB in sync, backup and restore takes long time

C is correct as DMS takes care both full load and ongoing replication, see this youtube video <https://www.youtube.com/watch?v=VhXDa9SPDLw>

D is right as from PostgreSQL to Amazon Aurora PostgreSQL you need AWS Schema Conversion Tool, see

<https://aws.amazon.com/dms/schema-conversion-tool/>

E is out monitor itself doesn't perform the replication work, if we have to choose 3 options then we can have E selected

upvoted 2 times

pentium75 1 year, 5 months ago

https://docs.aws.amazon.com/zh_cn/dms/latest/sbs/chap-manageddatabases.postgresql-rds-postgresql-ongoing-replication.html literally says that you must "configure the ongoing replication task"

upvoted 2 times

Mikado211 1 year, 6 months ago

Well technically when you operate such task, you must create a database on the cloud, then operate a migration using DMS and none of the propositions give you those two tasks separately. Sometimes those questions can be really frustrating.

upvoted 1 times

Amitabha09 1 year, 8 months ago

- C. Create an AWS Database Migration Service (AWS DMS) replication server.
- E. Create an Amazon EventBridge (Amazon CloudWatch Events) rule to monitor the database synchronization.

AWS DMS can replicate data from on-premises databases to Aurora PostgreSQL in real time, so the on-premises database will remain online and accessible during the migration. AWS DMS can also automatically convert the database schema, so there is no need to use AWS SCT.

An Amazon EventBridge rule can be used to monitor the database synchronization and send notifications if any errors occur. This is important because it allows the solutions architect to quickly identify and resolve any issues that may arise during the migration.

A database backup of the on-premises database is not necessary because AWS DMS will replicate the data in real time. Creating an ongoing replication task is not necessary because AWS DMS will automatically create an ongoing replication task when the replication server is created.

upvoted 2 times

David_Ang 1 year, 7 months ago

Mate you can monitor everything you want but it is not going to make sure the synchronization is working, an alert is not going to help.

upvoted 4 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: AC

Create an AWS Database Migration Service (AWS DMS) replication server then create an ongoing replication task

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #137

Topic 1

A company uses AWS Organizations to create dedicated AWS accounts for each business unit to manage each business unit's account independently upon request. The root email recipient missed a notification that was sent to the root user email address of one account. The company wants to ensure that all future notifications are not missed. Future notifications must be limited to account administrators.

Which solution will meet these requirements?

A. Configure the company's email server to forward notification email messages that are sent to the AWS account root user email address to all users in the organization.

B. Configure all AWS account root user email addresses as distribution lists that go to a few administrators who can respond to alerts. Configure AWS account alternate contacts in the AWS Organizations console or programmatically.


Most Voted

C. Configure all AWS account root user email messages to be sent to one administrator who is responsible for monitoring alerts and forwarding those alerts to the appropriate groups.

D. Configure all existing AWS accounts and all newly created accounts to use the same root user email address. Configure AWS account alternate contacts in the AWS Organizations console or programmatically.

Correct Answer: B

Community vote distribution



Comments

123jhl0 Highly Voted 2 years, 7 months ago

Selected Answer: B

Use a group email address for the management account's root user
https://docs.aws.amazon.com/organizations/latest/userguide/orgs_best-practices_mgmt-acct.html#best-practices_mgmt-acct_email-address

upvoted 28 times

cookieMr Highly Voted 1 year, 11 months ago

Selected Answer: B

Option B ensures that all future notifications are not missed by configuring the AWS account root user email addresses as distribution lists that are monitored by a few administrators. By setting up alternate contacts in the AWS Organizations console

notifications that are monitored by a few administrators by setting up alternate contacts in the AWS Organizations console or programmatically, the notifications can be sent to the appropriate administrators responsible for monitoring and responding to alerts. This solution allows for centralized management of notifications and ensures they are limited to account administrators.

- A. Floods all users with notifications, lacks granularity.
- C. Manual forwarding introduces delays, centralizes responsibility.
- D. No flexibility for specific account administrators, limits customization.

upvoted 14 times

15df3d0 Most Recent 3 months, 1 week ago

Selected Answer: D

[//docs.aws.amazon.com/organizations/latest/userguide/orgs_best-practices_mgmt-acct.html#best-practices_mgmt-acct_email-address](https://docs.aws.amazon.com/organizations/latest/userguide/orgs_best-practices_mgmt-acct.html#best-practices_mgmt-acct_email-address)

upvoted 1 times

PaulGa 8 months, 4 weeks ago

Selected Answer: B

Ans B - as opposed to option D, because the organisation account structure implies there is more than one root account: "The root email recipient missed a notification that was sent to the root user email address of one account."

upvoted 2 times

MatAlves 9 months, 2 weeks ago

Selected Answer: B

"Use a group email address for root user credentials:

Use an email address that is managed by your business and forwards received messages directly to a group of users. If AWS must contact the owner of the account, this approach reduces the risk of delays in responding, even if individuals are on vacation, out sick, or have left the business. The email address used for the root user should not be used for other purposes."

<https://docs.aws.amazon.com/IAM/latest/UserGuide/root-user-best-practices.html#ru-bp-group>

upvoted 2 times

MandAsh 1 year ago

Lol.. How is this AWS related question. Isnt it general knowledge.

upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

No idea why "D" would be correct answer unless there is some missing context in the question or the answer. "B" is best practice as pointed out in other links.

upvoted 1 times

David_Ang 1 year, 7 months ago

Selected Answer: B

the only answer with sense is "B", because "A" is not exclusive, "C" is exactly the case the want to avoid, and "D" just don't make sense

upvoted 2 times

tom_cruise 1 year, 8 months ago

Selected Answer: B

distribution list is the way to go

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

The reasons are:

Alternate contacts allow defining other users to receive root emails.

Distribution lists ensure multiple admins get notified.

Limits notifications to account admins rather than all users.

Using the same root email address for all accounts (Option D) is not recommended.

Relying on one admin or external forwarding (Options A, C) introduces delays or single points of failure.

Keying on one admin or external forwarding (Options A, C) introduces delays or single points of failure.

upvoted 2 times

Itsume 1 year, 12 months ago

all admins need access or else some wont get the right mails and cant do their job, sending it only to a few would disrupt the workflow so it is D

upvoted 1 times

fishy_resolver 2 years ago

Selected Answer: D

From the links provided below there are no mention of having a distribution list capability within AWS:

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_best-practices_mgmt-acct.html#best-practices_mgmt-acct_email-address

As per link for best practices:

Use a group email address for the management account's root user!

upvoted 2 times

Abrar2022 2 years ago

The clue is in the pudding!!

Question: account "administrators"

Answer: Configure all AWS account root user email addresses as distribution lists that go to a few "administrators"

upvoted 3 times

Rainchild 2 years, 1 month ago

Selected Answer: B

Option A: wrong - sends email to everybody

Option B: correct (but sub-optimal because distribution lists aren't all that secure)

Option C: wrong - single point of failure on the new administrator

Option D: wrong - each root email address must be unique, you can't change them all to the same one

upvoted 3 times

jdr75 2 years, 2 months ago

Selected Answer: B

The more aligned answer to this article:

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_best-practices_mgmt-acct.html#best-practices_mgmt-acct_email-address

is B.

D would be best if it'd said that the email you configure as "root user email address" will be a distribution list.

The phrase "all future notifications are not missed" points to D, cos' it said:

".. and all newly created accounts to use the same root user email address"

so the future account that will be created will be covered with the business policy.

It's not 100% clear, but I'll choose B.

upvoted 3 times

TheAbsoluteTruth 2 years, 2 months ago

Una pregunta si la gente va votando las preguntas por que los administradores no cambian la respuesta correcta. Es a interpretación y ya?

upvoted 1 times

jdr75 2 years, 2 months ago

El administrador de "examtopics" pasa olímpicamente de marcar la respuesta correcta y es evidente que muchas respuestas que indica como "correctas" no lo son. Dice muy poco del servicio que dan.

upvoted 1 times

jaswantn 2 years, 2 months ago

Using the method of crossing out the option that does not fit....

Option A: address to all users of organization (wrong)

Option A: address to all users of organization (wrong)

Option B: go to a few administration who can respond to alerts (question says to send notification to administrators not a selected few)

Option C: send to one administrator and giving him responsibility (wrong)

Option D: correct (as this is the one option left after checking all others).

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #138

Topic 1

A company runs its ecommerce application on AWS. Every new order is published as a message in a RabbitMQ queue that runs on an Amazon EC2 instance in a single Availability Zone. These messages are processed by a different application that runs on a separate EC2 instance. This application stores the details in a PostgreSQL database on another EC2 instance. All the EC2 instances are in the same Availability Zone.

The company needs to redesign its architecture to provide the highest availability with the least operational overhead. What should a solutions architect do to meet these requirements?

A. Migrate the queue to a redundant pair (active/standby) of RabbitMQ instances on Amazon MQ. Create a Multi-AZ Auto Scaling group for EC2 instances that host the application. Create another Multi-AZ Auto Scaling group for EC2 instances that host the PostgreSQL database.

B. Migrate the queue to a redundant pair (active/standby) of RabbitMQ instances on Amazon MQ. Create a Multi-AZ Auto Scaling group for EC2 instances that host the application. Migrate the database to run on a Multi-AZ deployment of Amazon RDS for PostgreSQL. **Most Voted**

C. Create a Multi-AZ Auto Scaling group for EC2 instances that host the RabbitMQ queue. Create another Multi-AZ Auto Scaling group for EC2 instances that host the application. Migrate the database to run on a Multi-AZ deployment of Amazon RDS for PostgreSQL.

D. Create a Multi-AZ Auto Scaling group for EC2 instances that host the RabbitMQ queue. Create another Multi-AZ Auto Scaling group for EC2 instances that host the application. Create a third Multi-AZ Auto Scaling group for EC2 instances that host the PostgreSQL database

Correct Answer: B

Community vote distribution

B (100%)

Comments

123jhl0 **Highly Voted** 2 years, 7 months ago

Selected Answer: B

Migrating to Amazon MQ reduces the overhead on the queue management. C and D are dismissed. Deciding between A and B means deciding to go for an AutoScaling group for EC2 or an RDS for PostgreSQL (both multi- AZ). The RDS option has less operational impact, as provide as a service the tools and software required. Consider for instance, the effort to add an additional node like a read replica, to the DB.
<https://docs.aws.amazon.com/amazon-mq/latest/developer-guide/active-standby-broker-deployment.html>

<https://docs.aws.amazon.com/amazon-mq/latest/developer-guide/active-standby-broker-deployment.html>
<https://aws.amazon.com/rds/postgresql/>
upvoted 29 times

EKA_CloudGod 2 years, 6 months ago

This also helps anyone in doubt; <https://docs.aws.amazon.com/amazon-mq/latest/developer-guide/active-standby-broker-deployment.html>
upvoted 2 times

UWSFish 2 years, 7 months ago

Yes but active/standby is fault tolerance, not HA. I would concede after thinking about it that B is probably the answer that will be marked correct but its not a great question.
upvoted 4 times

Buruguduystunstugudunstuy Highly Voted 2 years, 5 months ago

Selected Answer: B

To meet the requirements of providing the highest availability with the least operational overhead, the solutions architect should take the following actions:

* By migrating the queue to Amazon MQ, the architect can take advantage of the built-in high availability and failover capabilities of the service, which will help ensure that messages are delivered reliably and without interruption.

* By creating a Multi-AZ Auto Scaling group for the EC2 instances that host the application, the architect can ensure that the application is highly available and able to handle increased traffic without the need for manual intervention.

* By migrating the database to a Multi-AZ deployment of Amazon RDS for PostgreSQL, the architect can take advantage of the built-in high availability and failover capabilities of the service, which will help ensure that the database is always available and able to handle increased traffic.

Therefore, the correct answer is Option B.

upvoted 8 times

zx9r Most Recent 6 months, 2 weeks ago

Selected Answer: B

B makes the most sense

upvoted 1 times

jaradat02 10 months, 3 weeks ago

Selected Answer: B

B is the correct answer

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

CD, you cannot have EC2 scaling work with RabbitMQ as only once instance can be active

A: Is good but B is better

B: Correct due to usage of RDS for PG so less overhead

upvoted 3 times

Gizmo2022 6 months, 3 weeks ago

Thank you for explaining the answer

upvoted 1 times

chandu7024 1 year, 8 months ago

Agree with B

upvoted 1 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: B

B offers high availability and low operational overheads.

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

Option B is the best solution to meet the high availability and low overhead requirements:

Migrate the queue to redundant Amazon MQ
Use Auto Scaling groups across AZs for the application
Migrate the database to Multi-AZ RDS PostgreSQL
The reasons are:

Amazon MQ provides a managed, highly available RabbitMQ cluster
Multi-AZ Auto Scaling distributes the application across AZs
RDS PostgreSQL is managed, multi-AZ capable database
Together this architecture removes single points of failure
RDS and MQ reduce operational overhead over self-managed

upvoted 6 times

MNotABot 1 year, 11 months ago

B
least operational overhead (Amazon RDS for PostgreSQL --> hence AD out / C says EC2 so out --> Hence B)

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: B

Option B provides the highest availability with the least operational overhead. By migrating the queue to a redundant pair of RabbitMQ instances on Amazon MQ, the messaging system becomes highly available. Creating a Multi-AZ Auto Scaling group for EC2 instances hosting the application ensures that it can automatically scale and maintain availability across multiple Availability Zones. Migrating the database to a Multi-AZ deployment of Amazon RDS for PostgreSQL provides automatic failover and data replication across multiple Availability Zones, enhancing availability and reducing operational overhead.

A. Incorrect because it does not address the high availability requirement for the RabbitMQ queue and the PostgreSQL database.

C. Incorrect because it does not provide redundancy for the RabbitMQ queue and does not address the high availability requirement for the PostgreSQL database.

D. Incorrect because it does not address the high availability requirement for the RabbitMQ queue and does not provide redundancy for the application instances.

upvoted 4 times

Gary_Phillips_2007 2 years, 3 months ago

Selected Answer: B

B for me.

upvoted 1 times

techhb 2 years, 5 months ago

Selected Answer: B

B is right all explanations below are correct

upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: B

Option B is right answer

upvoted 1 times

Wpcorgan 2 years, 6 months ago

B for me

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #139

Topic 1

A reporting team receives files each day in an Amazon S3 bucket. The reporting team manually reviews and copies the files from this initial S3 bucket to an analysis S3 bucket each day at the same time to use with Amazon QuickSight. Additional teams are starting to send more files in larger sizes to the initial S3 bucket.

The reporting team wants to move the files automatically analysis S3 bucket as the files enter the initial S3 bucket. The reporting team also wants to use AWS Lambda functions to run pattern-matching code on the copied data. In addition, the reporting team wants to send the data files to a pipeline in Amazon SageMaker Pipelines.


What should a solutions architect do to meet these requirements with the LEAST operational overhead?

A. Create a Lambda function to copy the files to the analysis S3 bucket. Create an S3 event notification for the analysis S3 bucket. Configure Lambda and SageMaker Pipelines as destinations of the event notification. Configure s3:ObjectCreated:Put as the event type.

B. Create a Lambda function to copy the files to the analysis S3 bucket. Configure the analysis S3 bucket to send event notifications to Amazon EventBridge (Amazon CloudWatch Events). Configure an ObjectCreated rule in EventBridge (CloudWatch Events). Configure Lambda and SageMaker Pipelines as targets for the rule.

C. Configure S3 replication between the S3 buckets. Create an S3 event notification for the analysis S3 bucket. Configure Lambda and SageMaker Pipelines as destinations of the event notification. Configure s3:ObjectCreated:Put as the event type.

D. Configure S3 replication between the S3 buckets. Configure the analysis S3 bucket to send event notifications to Amazon EventBridge (Amazon CloudWatch Events). Configure an ObjectCreated rule in EventBridge (CloudWatch Events). Configure Lambda and SageMaker Pipelines as targets for the rule. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****Six_Fingered_Jose** **Highly Voted** 2 years, 7 months ago**Selected Answer: D**

i go for D here

A and B says you are copying the file to another bucket using lambda,
C and D just uses S3 replication to copy the files,

They are doing exactly the same thing while C and D do not require setting up of lambda, which should be more efficient

The question says the team is manually copying the files, automatically replicating the files should be the most efficient method vs manually copying or copying with lambda.

upvoted 31 times

vipyodha 1 year, 11 months ago

yes d because of least operational overhead and also s3 event notification can only send to sns.sqs.and lambda , not to sagemaker.eventbridge can send to sagemaker

upvoted 22 times

Tsige 7 months, 4 weeks ago

S3 Replication: Configuring S3 replication between the initial and analysis S3 buckets automates the process of moving files between the buckets without the need to manually copy files or run a Lambda function for this purpose. This reduces operational overhead.

S3 Event Notifications: Once files are replicated to the analysis bucket, you can configure S3 event notifications for the s3:ObjectCreated

event. This event triggers actions (such as invoking Lambda functions and sending data to SageMaker Pipelines) when new files are placed in the analysis bucket.

The answer is C

upvoted 3 times

Abdou1604 1 year, 8 months ago

but the reporting team also wants to use AWS Lambda functions to run pattern-matching code on the copied , S3 replication cons is copying everything

upvoted 2 times

pentium75 1 year, 5 months ago

The Lambda functions should run "on the copied data", so first copy, THEN run Lambda function, which is achieved by D.

upvoted 5 times

123jhl0 Highly Voted 2 years, 7 months ago

Selected Answer: B

C and D aren't answers as replicating the S3 bucket isn't efficient, as other teams are starting to use it to store larger docs not related to the reporting, making replication not useful.

As Amazon SageMaker Pipelines, ..., is now supported as a target for routing events in Amazon EventBridge, means the answer is B

<https://aws.amazon.com/about-aws/whats-new/2021/04/new-options-trigger-amazon-sagemaker-pipeline-executions/>

upvoted 18 times

JayBee65 2 years, 5 months ago

I think you are mis-interpreting the question. I think you need to use all files, including the ones provided by other teams, otherwise how can you tell what files to copy? I think the point of this statement is to show that more files are in use, and being copied at different times, rather than suggesting you need to differentiate between the two sources of files.

upvoted 9 times

KADSM 2 years, 7 months ago

Not sure how far lambda will cope up with larger files with the timelimit in place.

upvoted 4 times

byteb 1 year, 5 months ago

"The reporting team wants to move the files automatically to analysis S3 bucket as the files enter the initial S3 bucket." Replication is asynchronous, with lambda the data will be available faster. So I think A is the answer.

upvoted 1 times

vipyodha 1 year, 11 months ago

but B is not least operational overhead , D is least operational overhead

upvoted 2 times

idr75 2 years, 2 months ago

jan 7 2 years, 2 months ago

You misinterpret it ... the reporting team is overload, cos' more teams request their services uploading more data to the bucket. That's the reason reporting team need to automate the process. So ALL the bucket objects need to be copied to other bucket, and the replication is better an cheaper than use Lambda. So the answer is D.

upvoted 3 times

gianola Most Recent 1 week, 1 day ago

Selected Answer: D

No answer is valid as the question asked how to "move" the files and they are getting duplicated. Otherwise I'd go for D.

upvoted 1 times

Soliner_Bilgi_Teknolojileri 2 weeks, 3 days ago

Selected Answer: C

- ☐ S3 Replication: Replication job 100% automated
- ☐ S3 event notification: Triggered when copied files arrive
- ☐ Both Lambda and SageMaker Pipelines (via Step Function or Lambda proxy) can be given as destinations
- ☐ Minimum operational burden + full automation
- ☐ s3:ObjectCreated:Put → appropriate trigger type

upvoted 1 times

ChhatwaniB 1 month, 2 weeks ago

Selected Answer: D

Amazon S3 can send event notification messages to the following destinations.

Amazon Simple Notification Service (Amazon SNS) topics
Amazon Simple Queue Service (Amazon SQS) queues
AWS Lambda
Amazon EventBridge

However, only one destination type can be specified for each event notification.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/notification-how-to-event-types-and-destinations.html>

upvoted 2 times

SirDNS 2 months, 3 weeks ago

Selected Answer: C

Why do we need to add EventBridge (additional layer/operational overhead) when S3 Event Notification does the job?

upvoted 1 times

AwsAbhiKumar 3 months, 1 week ago

Selected Answer: C

The question is little confusing between C and D.

Option C suggest to use S3 events notification in combination with AWS lambda and SageMaker. But S3 events notification don't have native integration to directly trigger Amazon SageMaker Pipelines. But it can work around this limitation by having the S3 event trigger a Lambda function, and then that Lambda function can call the SageMaker Pipelines API to start a pipeline execution.(Option doesn't directly suggests this combo)

Option D is fine but it uses S3 replication combined with EventBridge for notifications, which introduces an extra layer (EventBridge rules) that is unnecessary since S3 event notifications can directly trigger Lambda and SageMaker Pipelines. as suggested above.

upvoted 2 times

hpirnaj 5 months, 1 week ago

Selected Answer: C

c is the answer , S3 event notification can directly invoke a Lambda function, which can then in turn trigger a SageMaker Pipeline execution, effectively allowing an S3 event to initiate a SageMaker pipeline through a Lambda intermediary. we dont need to invoke EventBridge

upvoted 3 times

salman7540 5 months, 3 weeks ago

Selected Answer: D

S3 replication can use to copy files in different buckets so we don't need lambda.

S3 events can't be sent directly to sagemaker so we have to utilise eventbridge who supports many targets including sagemaker.

upvoted 5 times

rmanuraj 5 months, 4 weeks ago

Selected Answer: D

In the case of S3 event notification only one destination type can be specified for each event notification.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/notification-how-to-event-types-and-destinations.html#supported-notification-destinations>

upvoted 3 times

PSH123 6 months, 2 weeks ago

Selected Answer: C

gpt said 'C' is solution

upvoted 1 times

PaulGa 8 months, 3 weeks ago

Selected Answer: D

Ans D - least operational overhead using replication; I was initially going for Ans C until I spotted S3 event notification can only send to SQS, SNS, Lambda - not directly to Sagemaker; but Eventbridge can send to Sagemaker. Not sure why author prefers A...?

upvoted 4 times

jatric 11 months, 1 week ago

Selected Answer: D

S3 event can't be use to notify sagemaker, So C can't be right option. AB required lambda which is not unnecessary

upvoted 4 times

lofzee 1 year ago

Answer is D because it requires least operational overhead and S3 replication does the copying for you.

Also read this <https://docs.aws.amazon.com/AmazonS3/latest/userguide/EventNotifications.html>

Lambda and Sagemaker are not supported destinations for S3 Event Notifications

upvoted 4 times

lofzee 1 year ago

Sorry i meant Sagemaker is not supported as an S3 Event Notification. Lambda is though. Still doesn't change what the answer is.... D

upvoted 4 times

andyngh86 1 year, 4 months ago

I go for C, because option C no need to configure event notifications, but D need to extra work to configure the event notification, for the least operation, option C is best choice

upvoted 1 times

Marco_St 1 year, 6 months ago

Selected Answer: D

B is the first option I denied. Since it makes the event happens inside the analysis bucket to trigger the lambda function. so if the lambda function is running code to copy files from initial bucket to analysis bucket. Then this lambda function should be triggered by the event in initial bucket like once the data reaches in the initial bucket then lambda is triggered. D is the answer.

upvoted 2 times

AntonioMinolfi 1 year, 8 months ago

Selected Answer: D

Utilizing a lambda function would introduce additional operational overhead, eliminating options A and B. S3 replication offers a simpler setup and efficiently accomplishes the task. S3 notifications cannot use SageMaker as a destination; the permissible destinations include SQS, SNS, Lambda, and Eventbridge, so C is out.

upvoted 11 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #140

Topic 1

A solutions architect needs to help a company optimize the cost of running an application on AWS. The application will use Amazon EC2 instances, AWS Fargate, and AWS Lambda for compute within the architecture.

The EC2 instances will run the data ingestion layer of the application. EC2 usage will be sporadic and unpredictable. Workloads that run on EC2 instances can be interrupted at any time. The application front end will run on Fargate, and Lambda will serve the API layer. The front-end utilization and API layer utilization will be predictable over the course of the next year.

Which combination of purchasing options will provide the MOST cost-effective solution for hosting this application? (Choose two.)

- A. Use Spot Instances for the data ingestion layer **Most Voted**
- B. Use On-Demand Instances for the data ingestion layer
- C. Purchase a 1-year Compute Savings Plan for the front end and API layer. **Most Voted**
- D. Purchase 1-year All Upfront Reserved instances for the data ingestion layer.
- E. Purchase a 1-year EC2 instance Savings Plan for the front end and API layer.

Correct Answer: AC*Community vote distribution*

AC (100%)

Comments**SimonPark** **Highly Voted** 2 years, 7 months ago**Selected Answer: AC**

EC2 instance Savings Plan saves 72% while Compute Savings Plans saves 66%. But according to link, it says "Compute Savings Plans provide the most flexibility and help to reduce your costs by up to 66%. These plans automatically apply to EC2 instance usage regardless of instance family, size, AZ, region, OS or tenancy, and also apply to Fargate and Lambda usage." EC2 instance Savings Plans are not applied to Fargate or Lambda

upvoted 21 times

aba2s **Highly Voted** 2 years, 5 months ago**Selected Answer: AC**

Compute Savings Plans can be used for EC2 instances and Fargate. Whereas EC2 Savings Plans support EC2 only.

upvoted 9 times

David **Most Recent** 2 months, 2 weeks ago

Pauliga most recent 8 months, 3 weeks ago

Selected Answer: AC

Ans A, C -

A: Spot obvious for unpredictable, 'don't care' usage

C: Not so obvious... but its more than just EC2 - its about Compute power using Fargate, Lambda, API call processing so it has to be C (as opposed to E)

upvoted 2 times

huaze_jei 9 months, 1 week ago

Selected Answer: AC

Be mindful that the question is asking about API. So it should be Compute Savings Plans.

If it is for EC2, the Reserved Instance will be correct.

upvoted 3 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: AC

Compute Savings Plans can also apply to Fargate and Lambda Usage.

upvoted 5 times

AKBM7829 1 year, 9 months ago

BC is the answer

data ingestion = Spot Instance but

Keyword "Usage Unpredictable" : On-Demand

and for APi its Compute Savings Plan

upvoted 1 times

awashenko 1 year, 8 months ago

Spot instances can auto scale so Spot instance is correct.

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: AC

The two most cost-effective purchasing options for this architecture are:

A) Use Spot Instances for the data ingestion layer

C) Purchase a 1-year Compute Savings Plan for the front end and API layer

The reasons are:

Spot Instances provide the greatest savings for flexible, interruptible EC2 workloads like data ingestion.

Savings Plans offer significant discounts for predictable usage like the front end and API layer.

All Upfront and partial/no Upfront RI's don't align well with the sporadic EC2 usage.

On-Demand is more expensive than Spot for flexible EC2 workloads.

By matching purchasing options to the workload patterns, Spot for unpredictable EC2 and Savings Plans for steady-state usage, the solutions architect optimizes cost efficiency.

upvoted 3 times

cookieMr 1 year, 11 months ago

Selected Answer: AC

Using Spot Instances for the data ingestion layer will provide the most cost-effective option for sporadic and unpredictable workloads, as Spot Instances offer significant cost savings compared to On-Demand Instances (Option A).

Purchasing a 1-year Compute Savings Plan for the front end and API layer will provide cost savings for predictable utilization over the course of a year (Option C).

Option B is less cost-effective as it suggests using On-Demand Instances for the data ingestion layer, which does not take advantage of cost-saving opportunities.

Option D suggests purchasing 1-year All Upfront Reserved instances for the data ingestion layer, which may not be optimal for sporadic and unpredictable workloads.

sporadic and unpredictable workloads.

Option E suggests purchasing a 1-year EC2 instance Savings Plan for the front end and API layer, but Compute Savings Plans are typically more suitable for predictable workloads.

upvoted 4 times

Abrar2022 2 years ago

Spot instances for data injection because the task can be terminated at anytime and tolerate disruption. Compute Saving Plan is cheaper than EC2 instance Savings plan.

upvoted 2 times

Abrar2022 2 years ago

EC2 instance Savings Plans are not applied to Fargate or Lambda

upvoted 1 times

Noviiiiice 2 years, 2 months ago

Why not B?

upvoted 1 times

SkyZeroZx 2 years, 2 months ago

because onDemand is more expensive than spot additionally that the workload has no problem with being interrupted at any time

upvoted 2 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: AC

To optimize the cost of running this application on AWS, you should consider the following options:

A. Use Spot Instances for the data ingestion layer

C. Purchase a 1-year Compute Savings Plan for the front-end and API layer

Therefore, the most cost-effective solution for hosting this application would be to use Spot Instances for the data ingestion layer and to purchase either a 1-year Compute Savings Plan or a 1-year EC2 instance Savings Plan for the front-end and API layer.

upvoted 3 times

AKBM7829 1 year, 9 months ago

Yes, but in the question it also states that it is 'Unpredictable' So, On-Demand is suitable over Spot Instance right which makes BC as the answer

upvoted 1 times

lofzee 1 year ago

the question clearly says "can be interrupted at any time" - anything that mentions these words with cost saving, you should automatically think Spot instances

upvoted 2 times

awashenko 1 year, 8 months ago

Spot instances can auto scale so Spot is still correct.

upvoted 2 times

techhb 2 years, 5 months ago

Selected Answer: AC

Too obvious answer.

upvoted 1 times

berks 2 years, 5 months ago

Selected Answer: AC

AC

can be interrupted at any time => spot

upvoted 3 times

TECHNOWARRIOR 2 years, 5 months ago

A,E::

Savings Plan — EC2

Savings Plan offers almost the same savings from a cost as RIs and adds additional Automation around how the savings are being applied. One way to understand is to say that EC2 Savings Plan are Standard Reserved Instances with automatic switching depending on Instance types being used within the same instance family and additionally applied to ECS Fargate and Lambda.

Savings Plan — Compute

Savings Plan offers almost the same savings from a cost as RIs and adds additional Automation around how the savings are being applied. For example, they provide flexibility around instance types and regions so that you don't have to monitor new instance types that are being launched. It is also applied to Lambda and ECS Fargate workloads. One way to understand is to say that Compute Savings Plan are Convertible Reserved Instances with automatic switching depending on Instance types being used.

upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: AC

A and C

upvoted 1 times

rjam 2 years, 7 months ago

its A and C . <https://www.densify.com/finops/aws-savings-plan>

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #141

Topic 1


A company runs a web-based portal that provides users with global breaking news, local alerts, and weather updates. The portal delivers each user a personalized view by using mixture of static and dynamic content. Content is served over HTTPS through an API server running on an Amazon EC2 instance behind an Application Load Balancer (ALB). The company wants the portal to provide this content to its users across the world as quickly as possible.

How should a solutions architect design the application to ensure the LEAST amount of latency for all users?

- A. Deploy the application stack in a single AWS Region. Use Amazon CloudFront to serve all static and dynamic content by specifying the ALB as an origin. **Most Voted**
- B. Deploy the application stack in two AWS Regions. Use an Amazon Route 53 latency routing policy to serve all content from the ALB in the closest Region.
- C. Deploy the application stack in a single AWS Region. Use Amazon CloudFront to serve the static content. Serve the dynamic content directly from the ALB.
- D. Deploy the application stack in two AWS Regions. Use an Amazon Route 53 geolocation routing policy to serve all content from the ALB in the closest Region.

Correct Answer: A

Community vote distribution



Comments

huiy **Highly Voted** 2 years, 8 months ago

Selected Answer: A

Answer is A.

Amazon CloudFront is a web service that speeds up distribution of your static and dynamic web content

<https://www.examtopycs.com/discussions/amazon/view/81081-exam-aws-certified-solutions-architect-associate-saa-c02/>
upvoted 34 times

MutiverseAgent 1 year, 11 months ago

Also, option B does not use CloudFront which means all the traffic will go through the internet; So, despite deploying resources in two regions and using the lowest latency point, that public internet connection might probably be slower than a connection through a private aws network as Cloudfront can use.

upvoted 4 times

Six_Fingered_Jose Highly Voted 2 years, 7 months ago

Selected Answer: B

Answer should be B,

CloudFront reduces latency if its only static content, which is not the case here.
For Dynamic content, CF cant cache the content so it sends the traffic through the AWS Network which does reduces latency, but it still has to travel through another region.

For the case with 2 region and Route 53 latency routing, Route 53 detects the nearest resource (with lowest latency) and routes the traffic there. Because the traffic does not have to travel to resources far away, it should have the least latency in this case here.

upvoted 16 times

lofzee 1 year ago

All it takes is for you to go to the Amazon Cloudfront webpage hosted by AWS where it tells you "Amazon CloudFront is a content delivery network (CDN) service that helps you distribute your static and dynamic content quickly and reliably with high speed".

Answer is A bro. Your answer is B literally makes no sense. Do some studying.

upvoted 9 times

Aamee 2 years, 6 months ago

Can you pls. provide a ref. link from where this info. got extracted?

upvoted 1 times

manuelemg2007 1 year, 10 months ago

this is link <https://aws.amazon.com/es/blogs/aws-spanish/cloudfront-para-la-distribucion-de-contenido-estatico-y-dinamico/>

upvoted 4 times

Abdou1604 1 year, 8 months ago

What about accross the word :)

upvoted 4 times

Onimole 2 years, 7 months ago

Cf works for both static and dynamic content

upvoted 12 times

network_enthusiast Most Recent 1 month, 4 weeks ago

Selected Answer: B

B is the correct answer

upvoted 1 times

tch 2 months, 3 weeks ago

Selected Answer: C

Dynamic content can also be served via CloudFront by configuring the WordPress website as an origin. Since dynamic content includes personalized content, you need to configure CloudFront to forward certain HTTP cookies and HTTP headers as part of a request to your custom origin server. CloudFront uses the forwarded cookie values as part of the key that identifies a unique object in its cache. To ensure that you maximize the caching efficiency, you should configure CloudFront to only forward those HTTP cookies and HTTP headers that really vary the content (not cookies that are only used on the client side or by third-party applications, for example, for web analytics).

upvoted 1 times

Dharmarajan 4 months, 1 week ago

Selected Answer: A

A is correct, because the dynamic content is also cached - which means first time the data is brought in, then subsequently it does not have to go to the source to get the same content.

There are many ways to set these combinations up, even with Route 53, but of the given options, "A" seems to fit the bill best.

upvoted 1 times

satyaamm 4 months, 1 week ago

Selected Answer: A

CloudFront is suitable for both static and dynamic content.

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: B

The best approach combines multi-region deployment for dynamic content and Amazon CloudFront for static content. However, none of the options explicitly describe this combination. Among the given options, Option B is the best because:

It uses two regions to reduce latency for dynamic content.

Route 53 latency routing ensures users are directed to the closest region.

For a more optimal solution (not listed in the options):

Deploy the application stack in multiple AWS Regions.

Use Amazon CloudFront to serve static content from edge locations.

Use Route 53 latency routing to direct users to the closest region for dynamic content.

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

A : CloudFront is a global CDN that caches static content at edge locations, reducing latency for users worldwide.

However, dynamic content served through CloudFront with a single ALB origin will still incur latency for users far from the single AWS Region.

This solution does not address latency for dynamic content effectively

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

B - Deploying in two regions reduces latency for users closer to those regions.

Route 53 latency routing ensures users are directed to the closest region.

However, this solution does not leverage a CDN for static content, which could further reduce latency.

Verdict: Better than Option A but not optimal for static content delivery.

upvoted 1 times

rmanuraj 5 months, 4 weeks ago

Selected Answer: A

Amazon cloud front is a better choice in terms of delivering both static and dynamic content globally. Also option B says deploy the application stack to only two regions. but the use case is to access the portal globally. Don't think Amazon Route 53 latency routing policy will have a bigger impact in terms of low latency.

upvoted 2 times

PaulGa 8 months, 3 weeks ago

Selected Answer: B

Ans B - "If your application is hosted in multiple AWS Regions, you can improve performance for your users by serving their requests from the AWS Region that provides the lowest latency" ...because it needs to be dynamic: "Latency between hosts on the internet can change over time as a result of changes in network connectivity and routing. Latency-based routing is based on latency measurements taken over a period of time, and the measurements reflect these changes. A request that is routed to the Oregon Region this week might be routed to the Singapore Region next week."

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-latency.html>

upvoted 2 times

ChymKuBoy 11 months, 4 weeks ago

Selected Answer: B

B for sure
upvoted 1 times

ManikRoy 1 year, 1 month ago

With Amazon CloudFront, your end users connections are terminated at CloudFront locations closer to them, which helps in reducing the overall round trip time required to establish a connection. This is irrespective of static a dynamic content.
upvoted 2 times

ManikRoy 1 year, 1 month ago

Selected Answer: A

You can still have improved performance by distributing the dynamic traffic through CDN instead of ALB. Refer below link. Also for other 2 options, using just 2 other regions for world wide distribution doesn't make much of a sense.
<https://aws.amazon.com/cloudfront/dynamic-content/>
upvoted 3 times

Uzbekistan 1 year, 2 months ago

Selected Answer: C

CloudFront for Static Content: By leveraging Amazon CloudFront, static content such as images, stylesheets, and scripts can be cached and distributed globally across a network of edge locations. This ensures that users receive static content from the nearest edge location, reducing latency and improving performance.

Serve Dynamic Content from ALB: Since dynamic content requires real-time processing and cannot be effectively cached at edge locations, serving dynamic content directly from the Application Load Balancer (ALB) is appropriate. The ALB can handle dynamic requests efficiently within the AWS Region where the application is deployed.
upvoted 3 times

eb7be10 1 year, 1 month ago

C was my choice for the reasons stated here. What am I missing?
upvoted 3 times

lofzee 1 year ago

How can you SERVE content from a load balancer?
Amazon Cloudfront is designed for static and dynamic content. Why would you pick any other option that isn't A?
upvoted 3 times

tch 2 months, 3 weeks ago

<https://docs.aws.amazon.com/whitepapers/latest/best-practices-wordpress/dynamic-content.html>
upvoted 1 times

Parul25 1 year, 4 months ago

CloudFront improves the performance, availability, and security of your dynamic content but not the latency as compared to Route 53 Latency Routing policy. Hence option B
<https://aws.amazon.com/cloudfront/dynamic-content/>
upvoted 2 times

Parul25 1 year, 4 months ago

I choose option B.

While CloudFront can accelerate content delivery by caching static content at edge locations, it may not be the most effective solution in this scenario. Since the portal delivers a mixture of static and dynamic content, leveraging Route 53 latency routing for dynamic content delivery ensures that users are directed to the nearest AWS Region hosting the dynamic content.
upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: A

"Least amount of latency for all users" "across the world" = CloudFront, thus B and D are out. Also, deploying the stack in "two regions" would benefit those two regions, but not users "across the world".

CloudFront can also cache dynamic content, thus A.

upvoted 9 times

Bennyseg 1 year, 6 months ago

Selected Answer: A

Answer is option A:

Earth Networks uses a CDN so that they can provide dynamic and personalized web based content quickly to their users with very low latency and high performing response times. Specifically, they need to be able to provide local information to the end user, in near real time, and need a CDN that allows them to adjust things like time to live, query strings, and cookie information so that they can pass all that information back to the origin to pull just what the user needs.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #142

Topic 1

A gaming company is designing a highly available architecture. The application runs on a modified Linux kernel and supports only UDP-based traffic. The company needs the front-end tier to provide the best possible user experience. That tier must have low latency, route traffic to the nearest edge location, and provide static IP addresses for entry into the application endpoints. What should a solutions architect do to meet these requirements?

- A. Configure Amazon Route 53 to forward requests to an Application Load Balancer. Use AWS Lambda for the application in AWS Application Auto Scaling.
- B. Configure Amazon CloudFront to forward requests to a Network Load Balancer. Use AWS Lambda for the application in an AWS Application Auto Scaling group.
- C. Configure AWS Global Accelerator to forward requests to a Network Load Balancer. Use Amazon EC2 instances for the application in an EC2 Auto Scaling group. **Most Voted**
- D. Configure Amazon API Gateway to forward requests to an Application Load Balancer. Use Amazon EC2 instances for the application in an EC2 Auto Scaling group.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**dokaedu** **Highly Voted** 2 years, 7 months ago

Correct Answer: C

AWS Global Accelerator and Amazon CloudFront are separate services that use the AWS global network and its edge locations around the world. CloudFront improves performance for both cacheable content (such as images and videos) and dynamic content (such as API acceleration and dynamic site delivery). Global Accelerator improves performance for a wide range of applications over TCP or UDP by proxying packets at the edge to applications running in one or more AWS Regions. Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover. Both services integrate with AWS Shield for DDoS protection.

upvoted 87 times

praveenas400 2 years, 5 months ago

Explained very well. ty

upvoted 2 times

iCma 2 years, 6 months ago

Thank you, your explanation helped me to better understand even the answer of question 29
upvoted 4 times

stepman 2 years, 6 months ago

On top of this, lambda would not be able to run application that is running on a modified Linux kernel. The answer is C .
upvoted 9 times

Buruguduystunstugudunstuy Highly Voted 2 years, 5 months ago

Selected Answer: C

The correct answer is Option C. To meet the requirements;

* AWS Global Accelerator is a service that routes traffic to the nearest edge location, providing low latency and static IP addresses for the front-end tier. It supports UDP-based traffic, which is required by the application.

* A Network Load Balancer is a layer 4 load balancer that can handle UDP traffic and provide static IP addresses for the application endpoints.

* An EC2 Auto Scaling group ensures that the required number of Amazon EC2 instances is available to meet the demand of the application. This will help the front-end tier to provide the best possible user experience.

Option A is not a valid solution because Amazon Route 53 does not support UDP traffic.

Option B is not a valid solution because Amazon CloudFront does not support UDP traffic.

Option D is not a valid solution because Amazon API Gateway does not support UDP traffic.

upvoted 10 times

Buruguduystunstugudunstuy 2 years, 5 months ago

My mistake, correction on Option A, it is the Application Load Balancers do not support UDP traffic. They are designed to load balance HTTP and HTTPS traffic, and they do not support other protocols such as UDP.

upvoted 5 times

Dharmarajan Most Recent 4 months, 1 week ago

Selected Answer: C

Global Accelerator rules the roost in this use case! gives a static IP, and does everything asked here.

upvoted 1 times

PaulGa 8 months, 3 weeks ago

Selected Answer: C

Ans C - hint: "That tier must have low latency, route traffic to the nearest edge location, and provide static IP addresses for entry into the application endpoints."

"AWS Global Accelerator... provides static IP addresses that provide a fixed entry point to your applications and eliminate the complexity of managing specific IP addresses for different AWS Regions and Availability Zones... [to] always routes user traffic to the optimal endpoint based on performance, reacting instantly to changes in application health, your user's location, and policies that you configure."

<https://aws.amazon.com/global-accelerator/faqs/>

upvoted 3 times

huaze_lei 9 months, 1 week ago

Selected Answer: C

Non HTTP/S, the answer is always Global Accelerator without doubts. GA serves like Cloudfront on providing low latency through edge locations, with the exception of handling different protocols.

upvoted 3 times

lofzee 1 year ago

Whenever I see this line : "and provide static IP addresses for entry into the application endpoints." - my brain automatically thinks Global Accelerator.

upvoted 1 times

sidharthwader 1 year, 3 months ago

If the situation demands for UDP or some protocols that are not at application level then it would be better to use Global Accelerator and here they need top notch performance hence using it with NLB would be the best answer. Cloud Front does not support UDP nor does it support use of NLB

upvoted 2 times

Murtadhaceit 1 year, 6 months ago

Selected Answer: C

UDP: NLB.

Static IP: Global Accelerator.

upvoted 3 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: C

UDP, static IP = Global Accelerator and Network Load Balancer

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

AWS Global Accelerator provides static IP addresses that serve as a fixed entry point to application endpoints. This allows optimal routing to the nearest edge location.

Using a Network Load Balancer (NLB) allows support for UDP traffic, as NLBs can handle TCP and UDP protocols.

The application runs on a modified Linux kernel, so using Amazon EC2 instances directly will provide the needed customization and low latency.

The EC2 instances can be auto scaled based on demand to provide high availability.

API Gateway and Application Load Balancer are more suited for HTTP/HTTPS and REST API type workloads. For a UDP gaming workload, Global Accelerator + NLB + EC2 is a better architectural fit.

upvoted 3 times

cookieMr 1 year, 11 months ago

Selected Answer: C

AWS Global Accelerator is designed to improve the availability and performance of applications by routing traffic through the AWS global network to the nearest edge locations, reducing latency. By configuring AWS Global Accelerator to forward requests to a Network Load Balancer, UDP-based traffic can be efficiently distributed across multiple EC2 instances in an Auto Scaling group. Using Amazon EC2 instances for the application allows for customization of the Linux kernel and support for UDP-based traffic. This solution provides static IP addresses for entry into the application endpoints, ensuring consistent access for users.

Option A suggests using AWS Lambda for the application, but Lambda is not suitable for long-running UDP-based applications and may not provide the required low latency.

Option B suggests using CloudFront, which is primarily designed for HTTP/HTTPS traffic and does not have native support for UDP-based traffic.

Option D suggests using API Gateway, which is primarily used for RESTful APIs and does not support UDP-based traffic.

upvoted 4 times

Abrar2022 2 years ago

aws global accelerator provides static IP addresses.

upvoted 1 times

Bmarodi 2 years ago

Selected Answer: C

My choice is option C, due to the followings: Amazon Global accelerator route the traffic to nearest edge locations, it supports UDP-based traffic, and it provides static ip addresses as well, hence C is right answer.

upvoted 2 times

bakamon 2 years, 2 months ago

Answer : C

CloudFront : Doesn't support static IP addresses

ALB : Doesn't support UDP

upvoted 1 times

Devsin2000 2 years, 3 months ago

C - <https://aws.amazon.com/global-accelerator/>
upvoted 1 times

SilentMilli 2 years, 5 months ago

Selected Answer: C

To meet the requirements of providing low latency, routing traffic to the nearest edge location, and providing static IP addresses for entry into the application endpoints, the best solution would be to use AWS Global Accelerator. This service routes traffic to the nearest edge location and provides static IP addresses for the application endpoints. The front-end tier should be configured with a Network Load Balancer, which can handle UDP-based traffic and provide high availability. Option C, "Configure AWS Global Accelerator to forward requests to a Network Load Balancer. Use Amazon EC2 instances for the application in an EC2 Auto Scaling group," is the correct answer.

upvoted 1 times

techhb 2 years, 5 months ago

Selected Answer: C

C is obvious choice here.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #143

Topic 1


A company wants to migrate its existing on-premises monolithic application to AWS. The company wants to keep as much of the front-end code and the backend code as possible. However, the company wants to break the application into smaller applications. A different team will manage each application. The company needs a highly scalable solution that minimizes operational overhead.

Which solution will meet these requirements?

- A. Host the application on AWS Lambda. Integrate the application with Amazon API Gateway.
- B. Host the application with AWS Amplify. Connect the application to an Amazon API Gateway API that is integrated with AWS Lambda.
- C. Host the application on Amazon EC2 instances. Set up an Application Load Balancer with EC2 instances in an Auto Scaling group as targets.
- D. Host the application on Amazon Elastic Container Service (Amazon ECS). Set up an Application Load Balancer with Amazon ECS as the target. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

Ken701 **Highly Voted** 2 years, 7 months ago

I think the answer here is "D" because usually when you see terms like "monolithic" the answer will likely refer to microservices.
upvoted 40 times

Bevemo **Highly Voted** 2 years, 7 months ago

Selected Answer: D

D is organic pattern, lift and shift, decompose to containers, first making most use of existing code, whilst new features can be added over time with lambda+api gw later.
A is leapfrog pattern. requiring refactoring all code up front.
upvoted 21 times

Mimine87 **Most Recent** 2 months, 1 week ago

Selected Answer: D

Monolith to Smaller Apps (Microservices-Ready)

ECS (or EKS) is ideal when breaking a monolithic app into smaller services (aka microservices).

Different teams can manage individual containers running different parts of the app.

upvoted 1 times

Dharmarajan 4 months, 1 week ago

Selected Answer: D

D because the operational overhead is the smallest among the given options. The company may do all that breaking up of functionalities and let teams manage the parts, but operationally for the site, hosting in Containers is the lowest maintenance. No ASG tuning, no ALB limitations and so on.

upvoted 1 times

FlyingHawk 4 months, 1 week ago

Selected Answer: D

The company wants to keep as much of the front-end code and the backend code as possible, so containization is less code changes than B which uses lambda.

upvoted 1 times

zdi561 4 months, 1 week ago

Selected Answer: C

Auto scaling group meets highly scalable requirement. D is not right 1. it is unknown if the app can be containerized , 2. and it maintains EC2 as C, so D has no operational advantage. Microservice is not equivalent to container.

upvoted 1 times

PaulGa 8 months ago

Selected Answer: D

Ans D - hint: "...break the application into smaller applications"

upvoted 3 times

PaulGa 8 months, 3 weeks ago

Selected Answer: D

Ans D - hint: "The company wants to keep as much of the front-end code and the backend code as possible. However, the company wants to break the application into smaller applications." Containerisation will help the company achieve a scaleable, more manageable solution.

upvoted 2 times

pedro_vieira 10 months, 3 weeks ago

For the folks suggesting Amplify: Have any of you actually shipped anything on Amplify? There are tons of adaptations needed to port a monolith to Amplify, specially around the backend that will need severe refactor.

Answer D allows for decomposing the application into different containers, enabling a distributed monolith.

upvoted 4 times

ManikRoy 1 year, 1 month ago

Selected Answer: B

Amazon API Gateway and Amplify both server less. Also you can import your code from GitHub in the amplify.

upvoted 2 times

ManikRoy 1 year, 1 month ago

Option D does not mention AWS Fargate which would cover the 'least operational overhead ' part.

upvoted 1 times

ManikRoy 1 year, 1 month ago

The company wants to keep much of its existing code. So the preferable solution is ECS. However the option D does not mention AWS Fargate which would cover the 'least operational overhead ' part.

upvoted 2 times

jbkrishna 1 year, 2 months ago

Different teams working means "microservices based architecture" so basically decoupling the application ..u can achieve this only by containerizing the app so answer is D

upvoted 4 times

NayeraB 1 year, 4 months ago

Selected Answer: B

B allows for a serverless architecture using AWS Lambda functions, which are highly scalable and require minimal operational overhead. AWS Amplify can help in managing the front-end code, while Amazon API Gateway integrated with AWS Lambda can handle the backend services.

D imo is not the best option in this scenario. While ECS can be a good choice for containerized workloads, it might introduce more operational overhead compared to a serverless solution like AWS Lambda and AWS Amplify.

upvoted 4 times

awsgeek75 1 year, 4 months ago

Selected Answer: D

I have a problem with this question.

"The company wants to keep as much of the front-end code and the backend code as possible"

So containerization is the solution here (D)? ABC don't make much sense so I will go with D but using containers for FE/BE code and configuring ALB for ECS (hopefully for frontend containers) is a pain in practice. Maybe this is worded in a bad way.

upvoted 5 times

vip2 1 year, 5 months ago

Selected Answer: B

Original state: monolithic with FE and BE code

Wanted state: seperate to mutilple components for diff. teams as Microservices

B is correct to decouple monolithic to microservices.

D still keep monolithic application in ECS.

upvoted 3 times

06042022 1 year, 5 months ago

IT is B. AWS amplify.

AWS Amplify will help seperate FE and BE. I agree with MM_Korvinus answer.

upvoted 2 times

JTruong 1 year, 5 months ago

Selected Answer: D

<https://aws.amazon.com/tutorials/break-monolith-app-microservices-ecs-docker-ec2/module-three/>

This page explained clearly why D is the correct answer

upvoted 3 times


Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #144

Topic 1

A company recently started using Amazon Aurora as the data store for its global ecommerce application. When large reports are run, developers report that the ecommerce application is performing poorly. After reviewing metrics in Amazon CloudWatch, a solutions architect finds that the ReadIOPS and CPU Utilization metrics are spiking when monthly reports run. What is the MOST cost-effective solution?

- A. Migrate the monthly reporting to Amazon Redshift.
- B. Migrate the monthly reporting to an Aurora Replica. **Most Voted**
- C. Migrate the Aurora database to a larger instance class.
- D. Increase the Provisioned IOPS on the Aurora instance.

Correct Answer: B*Community vote distribution***Comments****Buruguduystunstugudunstuy** **Highly Voted** 2 years, 5 months ago**Selected Answer: B**

Option B: Migrating the monthly reporting to an Aurora Replica may be the most cost-effective solution because it involves creating a read-only copy of the database that can be used specifically for running large reports without impacting the performance of the primary database. This solution allows the company to scale the read capacity of the database without incurring additional hardware or I/O costs.

upvoted 17 times

Buruguduystunstugudunstuy 2 years, 5 months ago

The incorrect solutions are:

Option A: Migrating the monthly reporting to Amazon Redshift may not be cost-effective because it involves creating a new data store and potentially significant data migration and ETL costs.

Option C: Migrating the Aurora database to a larger instance class may not be cost-effective because it involves changing the underlying hardware of the database and potentially incurring additional costs for the larger instance.

Option D: Increasing the Provisioned IOPS on the Aurora instance may not be cost-effective because it involves paying for additional I/O capacity that may not be necessary for other workloads on the database.

upvoted 11 times

Mikado211 Highly Voted 1 year, 6 months ago

Selected Answer: B

Report = Aurora replica

upvoted 7 times

Dharmarajan Most Recent 4 months, 1 week ago

Selected Answer: B

B. Doesn't say "Read Replica" but I guess it is implied - what other kind of replica is there? only an active-active standby, still that would solve the issue.

upvoted 1 times

PaulGa 8 months ago

Selected Answer: B

Ans B - Using Aurora Replica means you can work off a snapshot without impacting the master DB

upvoted 2 times

PaulGa 8 months, 3 weeks ago

Selected Answer: B

Ans B - making reports is temporary feature every month, so creating Aurora Replica is sufficient and efficient. Redshift is unnecessary (and means further overheads); C, D permanently upgrade capability which is not needed

upvoted 2 times

1166ae3 11 months, 3 weeks ago

Selected Answer: D

This approach focuses directly on mitigating the observed IOPS spikes, which are likely contributing to the performance degradation during heavy report processing, without introducing additional complexities or higher operational costs associated with other options.

upvoted 1 times

lofzee 1 year ago

Aye B.

upvoted 1 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: B

Migrate the monthly reporting to an Aurora Replica

upvoted 4 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

Aurora Replicas utilize the same storage as the primary instance so there is no additional storage cost.

Replicas can be created and destroyed easily to match reporting needs.

The primary Aurora instance size does not need to be changed, avoiding additional cost.

Workload is offloaded from the primary instance, improving its performance.

No major software/configuration changes needed compared to options like Redshift.

upvoted 4 times

cd93 1 year, 10 months ago

I don't understand why doubling everything (instances, network cost, maintenance effort, and especially storage) can be considered "cost-saving" for a simple monthly report...

An instance upgrade can very well be much cheaper. This question is very vague and does not provide enough information.

upvoted 2 times

cd93 1 year, 10 months ago

Silly me, I thought upgrading instance type includes storage upgrade (increase read iops) lol. The question pointed out that hard drive is also a limiting factor, so correct answer is B.

upvoted 3 times

cookieMr 1 year, 11 months ago

Selected Answer: B

B is correct because migrating the monthly reporting to an Aurora Replica can offload the reporting workload from the primary Aurora instance, reducing the impact on its performance during large reports. Using an Aurora Replica provides scalability and allows the replica to handle the read-intensive reporting queries, improving the overall performance of the ecommerce application.

A is wrong because migrating to Amazon Redshift introduces additional costs and complexity, and it may not be necessary to switch to a separate data warehousing service for this specific issue.

C is wrong because simply increasing the instance class of the Aurora database may not be the most cost-effective solution if the performance issue can be resolved by offloading the reporting workload to an Aurora Replica.

D is wrong because increasing the Provisioned IOPS alone may not address the issue of spikes in CPU Utilization during large reports, as it primarily focuses on storage performance rather than overall database performance.

upvoted 5 times

Abrar2022 2 years ago

By using an Aurora Replica for running large reports, the primary database will be relieved of the additional read load, improving performance for the ecommerce application.

upvoted 2 times

Bmarodi 2 years ago

Selected Answer: B

Option B is right answer.

upvoted 1 times

studynoplay 2 years, 1 month ago

Finally a question where there are no controversies

upvoted 3 times

SilentMilli 2 years, 5 months ago

Selected Answer: B

The most cost-effective solution for addressing high ReadIOPS and CPU utilization when running large reports would be to migrate the monthly reporting to an Aurora Replica. An Aurora Replica is a read-only copy of an Aurora database that is updated in real-time with the primary database. By using an Aurora Replica for running large reports, the primary database will be relieved of the additional read load, improving performance for the ecommerce application. Option B, "Migrate the monthly reporting to an Aurora Replica," is the correct answer.

upvoted 2 times

career360guru 2 years, 5 months ago

Selected Answer: B

B is the best option

upvoted 2 times

sanket1990 2 years, 6 months ago

Selected Answer: B

B is correct

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #145

Topic 1


A company hosts a website analytics application on a single Amazon EC2 On-Demand Instance. The analytics software is written in PHP and uses a MySQL database. The analytics software, the web server that provides PHP, and the database server are all hosted on the EC2 instance. The application is showing signs of performance degradation during busy times and is presenting 5xx errors. The company needs to make the application scale seamlessly.

Which solution will meet these requirements MOST cost-effectively?

- A. Migrate the database to an Amazon RDS for MySQL DB instance. Create an AMI of the web application. Use the AMI to launch a second EC2 On-Demand Instance. Use an Application Load Balancer to distribute the load to each EC2 instance.
- B. Migrate the database to an Amazon RDS for MySQL DB instance. Create an AMI of the web application. Use the AMI to launch a second EC2 On-Demand Instance. Use Amazon Route 53 weighted routing to distribute the load across the two EC2 instances.
- C. Migrate the database to an Amazon Aurora MySQL DB instance. Create an AWS Lambda function to stop the EC2 instance and change the instance type. Create an Amazon CloudWatch alarm to invoke the Lambda function when CPU utilization surpasses 75%.
- D. Migrate the database to an Amazon Aurora MySQL DB instance. Create an AMI of the web application. Apply the AMI to a launch template. Create an Auto Scaling group with the launch template. Configure the launch template to use a Spot Fleet. Attach an Application Load Balancer to the Auto Scaling group. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

Konb **Highly Voted** 2 years, 1 month ago

Selected Answer: D

I was tempted to pick A but then I realized there are two key requirements:

- scale seamlessly
- cost-effectively

None of A-C give seamless scalability. A and B are about adding second instance (which I assume does not match to "scale seamlessly"). C is about changing instance type.

Therefore D is only workable solution to the scalability requirement.

Therefore D is only workable solution to the scalability requirement
upvoted 21 times

hpirnaj 5 months, 1 week ago

Aurora is way more expensive than MySQL. but . I would still pick option D because it will scale seamlessly with ASG and Aurora . A is not scale at all
upvoted 1 times

pbpally 2 years, 1 month ago

Yup. Got me too. I picked A, saw D, and then reread the "scale seamlessly" part. D is correct.
upvoted 6 times

NayeraB 1 year, 4 months ago

But wouldn't RDS scale as well? Also Spot instances seems like a bit of a risky decision
upvoted 2 times

genny Highly Voted 2 years, 2 months ago

Selected Answer: A

I wouldn't run my website on spot instances. Spot instances might be terminated at any time, and since I need to run analytics application it's not an option for me. And using route 53 for load balancing of 2 instances is an overkill. I go with A.
upvoted 12 times

AZ_Master 1 year, 6 months ago

It is spot fleet - not spot instances. They can include On-Demand instances and can also maintain the target capacity automatically.

A Spot Fleet is a set of Spot Instances and optionally On-Demand Instances that is launched based on criteria that you specify. The Spot Fleet selects the Spot capacity pools that meet your needs and launches Spot Instances to meet the target capacity for the fleet. By default, Spot Fleets are set to maintain target capacity by launching replacement instances after Spot Instances in the fleet are terminated.

Ref: <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/spot-fleet.html>
upvoted 15 times

eb7be10 1 year, 1 month ago

While Spot Fleet is temporary, the Auto-Scaling Group builds it back again and again. D is the best answer.
upvoted 2 times

network_enthusiast Most Recent 1 month, 4 weeks ago

Selected Answer: D

D is the answer
upvoted 1 times

zdi561 4 months, 1 week ago

Selected Answer: D

Both ec2 and db scale better in D.
upvoted 1 times

satyaamm 4 months, 1 week ago

Selected Answer: D

Only D refers to seamlessly scale ... rest are about increasing the number of instances or the size of the instance.
upvoted 1 times

Rcosmos 5 months ago

Selected Answer: D

Conclusão:

A opção D é a solução mais econômica e eficaz porque combina o poder de escalabilidade automática do grupo de Auto Scaling, a economia das instâncias Spot, e a alta performance do Amazon Aurora MySQL. Essa abordagem atende aos requisitos de custo, desempenho e escalabilidade de forma abrangente.

requisitos de custo, desempenho e escalabilidade de forma abrangente.

upvoted 1 times

PaulGa 8 months ago

Selected Answer: D

Ans D - Like Konb (1 year, 5 months ago) I almost picked A, but for the same reasons of seamless scaling it has to be D

upvoted 2 times

carlossarmient 8 months, 2 weeks ago

D aswer is forcing you to use the new "feature" of aws. if you have a performance issue with the ddbb, front and back all in on your ec2 intances. the best way to solve that issue is move de ddbb to RDS and create a new ec2 intances. BUT we need to sell the aws unique service.

upvoted 2 times

PaulGa 8 months, 3 weeks ago

Selected Answer: D

Ans D - I was tempted with A
Ans D - I was tempted with A but then thought that's too obvious and scaling might be an issue... so looking at Spot Fleets "Fleets provide the following features and benefits, enabling you to maximize cost savings and optimize availability and performance when running applications on multiple EC2 instances" it has to be Ans D.

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Fleets.html>

upvoted 2 times

bignatov 9 months, 2 weeks ago

Selected Answer: A

Option A is the best balance between cost-effectiveness and scalability. It allows the application to scale horizontally with minimal changes while ensuring the database is managed and can scale independently, reducing the risk of performance degradation during peak times.

upvoted 1 times

bignatov 9 months, 2 weeks ago

i think it is not D, because amazon aurora is much expensive and also the spot instances can be interrupted, even they are most cost effective in the requirements is not mentioned that the workload is stateless and can be interrupted.

upvoted 2 times

maryam_sh 10 months, 1 week ago

Selected Answer: A

A, not D > Using Spot Instances with an Auto Scaling group adds complexity and risk to the infrastructure

upvoted 1 times

MatAlves 9 months, 2 weeks ago

It is not "spot instances", but "spot fleet".

A Spot Fleet is a set of Spot Instances and optionally On-Demand Instances that is launched based on criteria that you specify. The Spot Fleet selects the Spot capacity pools that meet your needs and launches Spot Instances to meet the target capacity for the fleet. By default, Spot Fleets are set to maintain target capacity by launching replacement instances after Spot Instances in the fleet are terminated.

Ref: <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/spot-fleet.html>

upvoted 1 times

jaradat02 10 months, 3 weeks ago

Selected Answer: D

I don't think D is the optimal solution. We certainly can find other solutions that are more cost effective and fulfill the same requirements, but among the provided options, I think D is the most reasonable.

upvoted 2 times

gt520490aws 11 months, 3 weeks ago

PHP....

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: D

I'll pick D because it sounds fun :)

upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: D

"Scale seamlessly", none of A-C include scaling at all.

upvoted 2 times

xdkonorek2 1 year, 7 months ago

Selected Answer: D

spot instance receives 2 minutes interruption notice, it should be enough for requests to finish, it's quite unusual for app to run longer requests

only option D allow for seamless scaling with autoscaling group

upvoted 2 times

BrijMohan08 1 year, 9 months ago

Selected Answer: B

Option B is a cost-effective choice that combines the benefits of database migration to RDS, horizontal scaling with EC2 instances, and control over traffic distribution with Route 53 weighted routing, making it the best solution for the given requirements.

upvoted 2 times

pentium75 1 year, 5 months ago

But there's no scaling at all in B.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #146


Topic 1

A company runs a stateless web application in production on a group of Amazon EC2 On-Demand Instances behind an Application Load Balancer. The application experiences heavy usage during an 8-hour period each business day. Application usage is moderate and steady overnight. Application usage is low during weekends. The company wants to minimize its EC2 costs without affecting the availability of the application. Which solution will meet these requirements?

- A. Use Spot Instances for the entire workload.
- B. Use Reserved Instances for the baseline level of usage. Use Spot instances for any additional capacity that the application needs. **Most Voted**
- C. Use On-Demand Instances for the baseline level of usage. Use Spot Instances for any additional capacity that the application needs.
- D. Use Dedicated Instances for the baseline level of usage. Use On-Demand Instances for any additional capacity that the application needs.

Correct Answer: B

Community vote distribution



Comments

rob74 **Highly Voted** 2 years, 7 months ago

Selected Answer: B

In the Question is mentioned that it has o Demand instances...so I think is more cheapest reserved and spot
upvoted 26 times

Qjb8m9h **Highly Voted** 2 years, 6 months ago

Answer is B: Reserved is cheaper than on demand the company has. And it's meet the availabilty (HA) requirement as to spot instance that can be disrupted at any time.

PRICING BELOW.

On-Demand: 0% There's no commitment from you. You pay the most with this option.

Reserved : 40%-60%1-year or 3-year commitment from you. You save money from that commitment.

Spot 50%-90% Ridiculously inexpensive because there's no commitment from the AWS side.

upvoted 13 times

3d54049 Most Recent 1 month ago

Selected Answer: B

Reserved and spot instances are cheaper than on-demand ones.

upvoted 1 times

manal001 3 months, 3 weeks ago

Selected Answer: C

couldn't decide between B and C, but C seems right. The thing with reserved instances is that you have to commit for 1 or 3 years, which is not mentioned in this use case.

upvoted 2 times

zdi561 4 months, 1 week ago

Selected Answer: C

Though ASG is not mentioned but it is implied. Mixing on-demand and spot is the best to save money. Reserved is expensive because you need to pay on the capacity all times.

upvoted 1 times

satyaamm 4 months, 1 week ago

Selected Answer: B

Since Reserved Instances are cheaper than On-Demand and also Spot Instances are best suited here for additional requirements.

upvoted 1 times

PaulGa 8 months ago

Selected Answer: B

Ans B - Reserved guarantees baseline level operation; Spot for peaks - its stateless

upvoted 2 times

jaradat02 10 months, 3 weeks ago

Selected Answer: B

B is the correct answer, using reserved instances is definitely more cost effective than using on-demand instances.

upvoted 2 times

ChymKuBoy 11 months, 4 weeks ago

Selected Answer: B

B for sure

upvoted 1 times

ManikRoy 1 year, 1 month ago

Selected Answer: B

Agree with others

upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: B

This is a bit unclear, but B seems the best option of the ones given.

Usage is either "heavy" (during the 8 hours), "moderate and steady" (overnight) or "low" (during weekends). So there is always SOME usage, which could be covered by a few Reserved Instances (which would be cheaper than On-Demand Instances).

A - "Spot instances for the entire workload", might 'affect the availability of the application'

B - Seems the best answer

C - More expensive than B

D - Dedicated instances aka dedicated hardware -> very expensive

upvoted 6 times

awsgeek75 1 year, 4 months ago

Agree, very little clarity between B and C but B makes more sense.

upvoted 3 times

HackPack 1 year, 5 months ago

I vote for C:

Please explain me if I am wrong:

If application experiences heavy usage during an 8-hour period each business day and all other time we don't need them? it mean than on-demand price will be only 33% from total cost so saving will be near 66%, more than reserved instances all other load we can cover by spot instances.

So why it not C?

upvoted 1 times

dungtrungpham 1 year, 4 months ago

You got it wrong.

You need the application all the time (24/7) because it says: "moderate and steady overnight, low usage at the weekend", not 8 hours a day

upvoted 4 times

VladanO 1 year, 6 months ago

Selected Answer: C

On-Demand Instances are more appropriate than Reserved Instances because "The application is used heavily for a period of 8 hours every weekday" requirements.

upvoted 2 times

rcptryk 1 year, 6 months ago

Selected Answer: C

The answer should be C. Because if reserved is chosen, you have to pay for every hour. I calculate from this pages (if I'm wrong please correct me) <https://aws.amazon.com/ec2/pricing/reserved-instances/pricing/#:~:text=Reserved%20Instances%20provide%20you%20with,instances%20when%20you%20need%20them>.

Example: for t4g.nano

Reserved instances $(0.003 \times 24 \times 365) + (1.90 \times 12) = 49.08$

On demand instance $(0.0042 \times 8 \times 365) = 12.264$

it will be added spot instances

upvoted 3 times

pentium75 1 year, 5 months ago

"The baseline level of usage" is the minimum usage that is always there (even at night and during weekends), for THAT you can use Reserved Instance.

upvoted 4 times

Marco_St 1 year, 6 months ago

Selected Answer: B

B, since the application needs to be on 24/7 for business days; on weekends, it can be off at any moment. The question mentions something like 8 hour per business day but!!! this is just for heavy usage, the application is also on during overnight.

upvoted 2 times

Juliez 1 year, 7 months ago

Why it's not A ? the application is "stateless" so it can be interrupted at any moment and the spot option is the cheaper one.

upvoted 2 times

vi24 1 year, 2 months ago

The statelessness of a web application does not necessarily mean that it's okay to be interrupted. Statelessness refers to how the application handles requests and manages session data, not its ability to handle interruptions.

upvoted 1 times

pentium75 1 year, 5 months ago

But there might not be any Spot Instances available and the app would go offline.

upvoted 2 times

StudyAllNite 1 year, 7 months ago

Selected Answer: C

If we assume moderate usage of 8 hours on average every day a week, this should be on demand, since it is not a 24/7 server. There is downtime on the weekends and after the initial 8 hours.

upvoted 2 times

SVDK 1 year, 5 months ago

There is no downtime. The application runs all the time (even weekends). Weekends is the base workload which we cover with reserved instances, the higher workloads during the week is covered by spot instances.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #147

Topic 1

A company needs to retain application log files for a critical application for 10 years. The application team regularly accesses logs from the past month for troubleshooting, but logs older than 1 month are rarely accessed. The application generates more than 10 TB of logs per month.

Which storage option meets these requirements MOST cost-effectively?

A. Store the logs in Amazon S3. Use AWS Backup to move logs more than 1 month old to S3 Glacier Deep Archive.

B. Store the logs in Amazon S3. Use S3 Lifecycle policies to move logs more than 1 month old to S3 Glacier Deep Archive.

Most Voted

C. Store the logs in Amazon CloudWatch Logs. Use AWS Backup to move logs more than 1 month old to S3 Glacier Deep Archive.

D. Store the logs in Amazon CloudWatch Logs. Use Amazon S3 Lifecycle policies to move logs more than 1 month old to S3 Glacier Deep Archive.

Correct Answer: B

Community vote distribution

B (100%)

Comments

rjam **Highly Voted** 2 years, 6 months ago

Selected Answer: B

Why not AwsBackup? No Glacier Deep is supported by AWS Backup

<https://docs.aws.amazon.com/aws-backup/latest/devguide/s3-backups.html>

AWS Backup allows you to backup your S3 data stored in the following S3 Storage Classes:

- S3 Standard
- S3 Standard - Infrequently Access (IA)
- S3 One Zone-IA
- S3 Glacier Instant Retrieval
- S3 Intelligent-Tiering (S3 INT)

upvoted 11 times

tdkcumberland 2 years, 6 months ago

AWS BackUp costs something, setting up S3 LCP doesn't.

upvoted 8 times

cookieMr Highly Voted 1 year, 11 months ago

Selected Answer: B

B is the most cost-effective solution. Storing the logs in S3 and using S3 Lifecycle policies to transition logs older than 1 month to S3 Glacier Deep Archive allows for cost optimization based on data access patterns. Since logs older than 1 month are rarely accessed, moving them to S3 Glacier Deep Archive helps minimize storage costs while still retaining the logs for the required 10-year period.

A is incorrect because using AWS Backup to move logs to S3 Glacier Deep Archive can incur additional costs and complexity compared to using S3 Lifecycle policies directly.

C adds unnecessary complexity and costs by involving CloudWatch Logs and AWS Backup when direct management through S3 is sufficient.

D is incorrect because using S3 Lifecycle policies to move logs from CloudWatch Logs to S3 Glacier Deep Archive is not a valid option. CloudWatch Logs and S3 have separate storage mechanisms, and S3 Lifecycle policies cannot be applied directly to CloudWatch Logs.

upvoted 7 times

satyaamm Most Recent 4 months, 1 week ago

Selected Answer: B

S3 and S3 Lifecycle policies are the most suited here.

upvoted 1 times

PaulGa 8 months ago

Selected Answer: B

Ans B - S3 with Glacier Deep plus Lifecycle management

upvoted 2 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: B

S3 Lifecycle policies to the rescue

upvoted 3 times

Mamiololo 2 years, 5 months ago

B is correct..

upvoted 1 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: B

Option B (Store the logs in Amazon S3. Use S3 Lifecycle policies to move logs more than 1-month-old to S3 Glacier Deep Archive) would meet these requirements in the most cost-effective manner.

This solution would allow the application team to quickly access the logs from the past month for troubleshooting, while also providing a cost-effective storage solution for the logs that are rarely accessed and need to be retained for 10 years.

upvoted 2 times

career360guru 2 years, 5 months ago

Selected Answer: B

Option B is most cost effective. Moving logs to Cloudwatch logs may incur additional cost.

upvoted 2 times

Wpcorgan 2 years, 6 months ago

B is correct

upvoted 1 times

ArielSchivo 2 years, 7 months ago

Selected Answer: B

S3 + Glacier is the most cost effective

S3 + Glacier is the most cost effective.

upvoted 4 times

Bevemo 2 years, 7 months ago

Selected Answer: B

D works, archive cloudwatch logs to S3 but is an additional service to pay for over B.

upvoted 2 times

Aamee 2 years, 6 months ago

CloudWatch logs can't store around 10 TB of data per month I believe so both C and D options are ruled out already.

upvoted 1 times

masetromain 2 years, 7 months ago

Selected Answer: B

<https://www.examtopics.com/discussions/amazon/view/80772-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #148

Topic 1

A company has a data ingestion workflow that includes the following components:

An Amazon Simple Notification Service (Amazon SNS) topic that receives notifications about new data deliveries

An AWS Lambda function that processes and stores the data


The ingestion workflow occasionally fails because of network connectivity issues. When failure occurs, the corresponding data is not ingested unless the company manually reruns the job.

What should a solutions architect do to ensure that all notifications are eventually processed?

- A. Configure the Lambda function for deployment across multiple Availability Zones.
- B. Modify the Lambda function's configuration to increase the CPU and memory allocations for the function.
- C. Configure the SNS topic's retry strategy to increase both the number of retries and the wait time between retries.
- D. Configure an Amazon Simple Queue Service (Amazon SQS) queue as the on-failure destination. Modify the Lambda function to process messages in the queue. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

bunnychip **Highly Voted** 2 years, 7 months ago

Selected Answer: D

ensure that all notifications are eventually processed

upvoted 17 times

Help2023 **Highly Voted** 2 years, 3 months ago

Selected Answer: D

This is why <https://docs.aws.amazon.com/sns/latest/dg/sns-message-delivery-retries.html>

upvoted 6 times

satyaamm **Most Recent** 4 months, 1 week ago

Selected Answer: D

SQS queues are suited here for message buffering.

upvoted 1 times

PaulGa 8 months ago

Selected Answer: D

Ans D - using SQS ensure the data is captured and not lost for processing

upvoted 2 times

ChymKuBoy 11 months, 4 weeks ago

Selected Answer: D

D for sure

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: D

Configure an Amazon Simple Queue Service (Amazon SQS) queue as the on-failure destination. Modify the Lambda function to process messages in the queue.

upvoted 5 times

CaoMengde09 2 years, 4 months ago

C is not the right answer since after several retries SNS discard the message which doesn't align with the requirement. D is the right answer

upvoted 4 times

CaoMengde09 2 years, 4 months ago

Best solution to process failed SNS notifications is using sns-dead-letter-queues (SQS Queue for reprocessing)

<https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html>

upvoted 4 times

SilentMilli 2 years, 5 months ago

Selected Answer: D

To ensure that all notifications are eventually processed, the best solution would be to configure an Amazon Simple Queue Service (SQS) queue as the on-failure destination for the SNS topic. This will allow the notifications to be retried until they are successfully processed. The Lambda function can then be modified to process messages in the queue, ensuring that all notifications are eventually processed. Option D, "Configure an Amazon Simple Queue Service (Amazon SQS) queue as the on-failure destination. Modify the Lambda function to process messages in the queue," is the correct answer.

upvoted 4 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: D

I choose Option D as the correct answer.

To ensure that all notifications are eventually processed, the solutions architect can set up an Amazon SQS queue as the on-failure destination for the Amazon SNS topic. This way, when the Lambda function fails due to network connectivity issues, the notification will be sent to the queue instead of being lost. The Lambda function can then be modified to process messages in the queue, ensuring that all notifications are eventually processed.

upvoted 5 times

techhb 2 years, 5 months ago

Selected Answer: D

Option D to ensure that all notifications are eventually processed you need to use SQS.

upvoted 3 times

career360guru 2 years, 5 months ago

Selected Answer: C

Option C is right option.

SNS does not have any "On Failure" delivery destination. One need to configure dead-letter queue and configure SQS to read from there. So given this option D is incorrect.

upvoted 2 times

JavBee65 2 years, 5 months ago

I don't think that's right

"A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully. Messages that can't be delivered due to client errors or server errors are held in the dead-letter queue for further analysis or reprocessing" from <https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html>.

This is pretty much what is being described in D.

Plus C will only retry message processing, and network problems could still prevent the message from being processed, but the question states "ensure that all notifications are eventually processed". So C does not meet the requirements but D does look to do this.

upvoted 8 times

Gajendr 1 year, 5 months ago

+ <https://docs.aws.amazon.com/sns/latest/dg/sns-message-delivery-retries.html>

""To keep the message after the retries specified in the delivery policy are exhausted, configure your subscription to move undeliverables messages to a dead-letter queue (DLQ). For more information"" So D

upvoted 1 times

NikaCZ 2 years, 5 months ago

Selected Answer: D

Is correct.

upvoted 1 times

NikaCZ 2 years, 5 months ago

If you want to ensure that all notifications are eventually processed you need to use SQS.

upvoted 2 times

Wajif 2 years, 6 months ago

Selected Answer: D

C isnt specific. Hence D

upvoted 1 times

LeGloupier 2 years, 6 months ago

Selected Answer: C

"on-failure destination" doesn't exist, only dead letter queue exist.

that's why I am leaning for C

upvoted 1 times

Wajif 2 years, 6 months ago

Dead letter queue doesnt exist in SNS. They are specifically saying a new queue will be configured for failures from SNS.

Hence D

upvoted 3 times

Wpcorgan 2 years, 6 months ago

D is correct

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #149

Topic 1

A company has a service that produces event data. The company wants to use AWS to process the event data as it is received. The data is written in a specific order that must be maintained throughout processing. The company wants to implement a solution that minimizes operational overhead. How should a solutions architect accomplish this?

- A. Create an Amazon Simple Queue Service (Amazon SQS) FIFO queue to hold messages. Set up an AWS Lambda function to process messages from the queue. **Most Voted**
- B. Create an Amazon Simple Notification Service (Amazon SNS) topic to deliver notifications containing payloads to process. Configure an AWS Lambda function as a subscriber.
- C. Create an Amazon Simple Queue Service (Amazon SQS) standard queue to hold messages. Set up an AWS Lambda function to process messages from the queue independently.
- D. Create an Amazon Simple Notification Service (Amazon SNS) topic to deliver notifications containing payloads to process. Configure an Amazon Simple Queue Service (Amazon SQS) queue as a subscriber.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: A

The correct solution is Option A. Creating an Amazon Simple Queue Service (Amazon SQS) FIFO queue to hold messages and setting up an AWS Lambda function to process messages from the queue will ensure that the event data is processed in the correct order and minimize operational overhead.

Option B is incorrect because using Amazon Simple Notification Service (Amazon SNS) does not guarantee the order in which messages are delivered.

Option C is incorrect because using an Amazon SQS standard queue does not guarantee the order in which messages are processed.

Option D is incorrect because using an Amazon SQS queue as a subscriber to an Amazon SNS topic does not guarantee the order in which messages are processed.

upvoted 6 times

cookieMr Highly Voted 1 year, 11 months ago

A is the correct solution. By creating an Amazon Simple Queue Service (Amazon SQS) FIFO queue to hold messages and setting up an AWS Lambda function to process messages from the queue, the company can ensure that the order of the event data is maintained throughout processing. SQS FIFO queues guarantee the order of messages and are suitable for scenarios where strict message ordering is required.

B is incorrect because Amazon Simple Notification Service (Amazon SNS) topics are not designed to preserve message order. SNS is a publish-subscribe messaging service and does not guarantee the order of message delivery.

C is incorrect because using an SQS standard queue does not guarantee the order of message processing. SQS standard queues provide high throughput and scale, but they do not guarantee strict message ordering.

D is incorrect because configuring an SQS queue as a subscriber to an SNS topic does not ensure message ordering. SNS topics distribute messages to subscribers independently, and the order of message processing is not guaranteed.

upvoted 5 times

Dharmarajan Most Recent 4 months, 1 week ago

Selected Answer: A

Bst wouldbe to set up a SQS FIFO queue! anytime "Order" is important, FIFO rules. Other options dont have ordering as a priority and they dont guarantee order.

upvoted 1 times

satyaammm 4 months, 1 week ago

Selected Answer: A

SQS FIFO queues maintain the order and hence are suitable here.

upvoted 1 times

PaulGa 8 months ago

Selected Answer: A

Ans A - SQS (FIFO) ensures data is processed in the order it is received

upvoted 2 times

jaradat02 10 months, 3 weeks ago

Selected Answer: A

"The data is written in a specific order that must be maintained throughout processing."

A without reading other options.

upvoted 2 times

theochan 1 year, 5 months ago

Selected Answer: A

Easiest question ever?

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: A

"specific order" = must be FIFO queue = only mentioned in A

upvoted 4 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: A

Create an Amazon Simple Queue Service (Amazon SQS) FIFO queue to hold messages. Set up an AWS Lambda function to process messages from the queue.

upvoted 2 times

cheese929 2 years, 1 month ago

Selected Answer: A

A is correct. Use FIFO to process in the specific order required

upvoted 3 times

WhericanIstart 2 years, 3 months ago

Selected Answer: A

Option A is correct...data is processed in the correct order
upvoted 2 times

techhb 2 years, 5 months ago

Only A is right option here.
upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: A

Option A is the best option.
upvoted 2 times

allect096 2 years, 5 months ago

Selected Answer: A

"The data is written in a specific order that must be maintained throughout processing" --> FIFO
upvoted 5 times

NikaCZ 2 years, 5 months ago

Selected Answer: A

specific order = FIFO
upvoted 2 times

k1kavi1 2 years, 6 months ago

Selected Answer: A

A is correct
upvoted 1 times

david76x 2 years, 6 months ago

Selected Answer: A

Definitely A
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #150

Topic 1


A company is migrating an application from on-premises servers to Amazon EC2 instances. As part of the migration design requirements, a solutions architect must implement infrastructure metric alarms. The company does not need to take action if CPU utilization increases to more than 50% for a short burst of time. However, if the CPU utilization increases to more than 50% and read IOPS on the disk are high at the same time, the company needs to act as soon as possible. The solutions architect also must reduce false alarms.

What should the solutions architect do to meet these requirements?

- A. Create Amazon CloudWatch composite alarms where possible. **Most Voted**
- B. Create Amazon CloudWatch dashboards to visualize the metrics and react to issues quickly.
- C. Create Amazon CloudWatch Synthetics canaries to monitor the application and raise an alarm.
- D. Create single Amazon CloudWatch metric alarms with multiple metric thresholds where possible.

Correct Answer: A

Community vote distribution



Comments

123jhl0 **Highly Voted** 2 years, 7 months ago

Selected Answer: A

Composite alarms determine their states by monitoring the states of other alarms. You can **use composite alarms to reduce alarm noise**. For example, you can create a composite alarm where the underlying metric alarms go into ALARM when they meet specific conditions. You then can set up your composite alarm to go into ALARM and send you notifications when the underlying metric alarms go into ALARM by configuring the underlying metric alarms never to take actions. Currently, composite alarms can take the following actions:

https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/Create_Composite_Alarm.html

upvoted 32 times

cookieMr **Highly Voted** 1 year, 11 months ago

Selected Answer: A

By creating composite alarms in CloudWatch, the solutions architect can combine multiple metrics, such as CPU utilization and read IOPS, into a single alarm. This allows the company to take action only when both conditions are met, reducing false alarms and focusing on meaningful alerts.

B can help in monitoring the overall health and performance of the application. However, it does not directly address the

B can help in monitoring the overall health and performance of the application. However, it does not directly address the specific requirement of triggering an action when CPU utilization and read IOPS exceed certain thresholds simultaneously.

C. Creating CloudWatch Synthetics canaries is useful for actively monitoring the application's behavior and availability. However, it does not directly address the specific requirement of monitoring CPU utilization and read IOPS to trigger an action.

D. Creating single CloudWatch metric alarms with multiple metric thresholds where possible can be an option, but it does not address the requirement of triggering an action only when both CPU utilization and read IOPS exceed their respective thresholds simultaneously.

upvoted 10 times

bishtr3 Most Recent 3 months, 1 week ago

Selected Answer: A

Rule : ALARM(CPUUtilizationTooHigh) AND ALARM(DiskReadOpsTooHigh)

upvoted 1 times

Dharmarajan 4 months, 1 week ago

Selected Answer: A

Composite metric is required. multiple occurrences of the threshold breach can be configured with Cloudwatch.

upvoted 1 times

zdi561 4 months, 1 week ago

Selected Answer: D

A is not right because you do not need multiple composite alarms. You need a single one. Both A and D work if they are single

upvoted 1 times

PaulGa 8 months ago

Selected Answer: A

Ans A - Cloudwatch composite alarms

upvoted 2 times

huaze_lei 9 months, 1 week ago

Selected Answer: A

With CloudWatch, you can combine several alarms into one composite alarm to create a summarized, aggregated health indicator over a whole application or group of resources. Composite alarms are alarms that determine their state by monitoring the states of other alarms. You define rules to combine the status of those monitored alarms using Boolean logic.

upvoted 2 times

NikuWithPakiya 1 year, 2 months ago

composite alarms are suited for scenarios where we have to combine the alarms for different metrics or dimensions, rather than for multiple threshold of the same metric. It contradict with option A.

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

Composite for multiple conditions like AND/OR combinations

B: This option just made me laugh. Lol, will someone just sit and look at this dashboard?

C: CW Synthetics canaries if for API

D: Single won't monitor multiple metrics

upvoted 3 times

Modulopi 1 year, 8 months ago

Selected Answer: A

A: Composite alarms determine their states by monitoring the states of other alarms. You can use composite alarms to reduce alarm noise. For example, you can create a composite alarm where the underlying metric alarms go into ALARM when they meet specific conditions.

upvoted 2 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: A

Composite alarms was designed to handle this scenario.

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: A

The key reasons are:

Composite alarms allow defining alarms with multiple metrics and conditions, like high CPU AND high read IOPS in this case. Composite alarms can avoid false positives triggered by a single metric spike.

Dashboards help visualize but won't take automated action. Synthetics tests application availability but doesn't address the metrics.

Single metric alarms with multiple thresholds can't correlate across metrics and may still trigger false positives.

Composite alarms allow acting quickly when both CPU and IOPS are high, per the stated need.

upvoted 5 times

Abrar2022 2 years ago

The composite alarm goes into ALARM state only if all conditions of the rule are met.

upvoted 3 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: A

Option A, creating Amazon CloudWatch composite alarms, is correct because it allows the solutions architect to create an alarm that is triggered only when both CPU utilization is above 50% and read IOPS on the disk are high at the same time. This meets the requirement to act as soon as possible if both conditions are met, while also reducing the number of false alarms by ensuring that the alarm is triggered only when both conditions are met.

upvoted 6 times

Buruguduystunstugudunstuy 2 years, 5 months ago

The incorrect solutions are:

In contrast, Option B, creating Amazon CloudWatch dashboards, would not directly address the requirement to trigger an alarm when both CPU utilization is high and read IOPS on the disk are high at the same time. Dashboards can be useful for visualizing metric data and identifying trends, but they do not have the capability to trigger alarms based on multiple metric thresholds.

Option C, using Amazon CloudWatch Synthetics canaries, may not be the best choice for this scenario, as canaries are used for synthetic testing rather than for monitoring live traffic. Canaries can be useful for monitoring the availability and performance of an application, but they may not be the most effective way to monitor the specific metric thresholds and conditions described in this scenario.

upvoted 4 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option D, creating single Amazon CloudWatch metric alarms with multiple metric thresholds, would not allow the solutions architect to create an alarm that is triggered only when both CPU utilization and read IOPS on the disk are high at the same time. Instead, the alarm would be triggered whenever any of the specified metric thresholds are exceeded, which may result in a higher number of false alarms.

upvoted 6 times

BENICE 2 years, 5 months ago

A is correct answer

upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: A

Option A

upvoted 1 times

Qjb8m9h 2 years, 6 months ago

The AWS::CloudWatch::CompositeAlarm type creates or updates a composite alarm. When you create a composite alarm, you specify a rule expression for the alarm that takes into account the alarm states of other alarms that you have created. The composite alarm goes into ALARM state only if all conditions of the rule are met.

The alarms specified in a composite alarm's rule expression can include metric alarms and other composite alarms.Using composite alarms can reduce alarm noise.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #151

Topic 1

A company wants to migrate its on-premises data center to AWS. According to the company's compliance requirements, the company can use only the ap-northeast-3 Region. Company administrators are not permitted to connect VPCs to the internet. Which solutions will meet these requirements? (Choose two.)

A. Use AWS Control Tower to implement data residency guardrails to deny internet access and deny access to all AWS Regions except ap-northeast-3. **Most Voted**

B. Use rules in AWS WAF to prevent internet access. Deny access to all AWS Regions except ap-northeast-3 in the AWS account settings.


C. Use AWS Organizations to configure service control policies (SCPS) that prevent VPCs from gaining internet access. Deny access to all AWS Regions except ap-northeast-3. **Most Voted**

D. Create an outbound rule for the network ACL in each VPC to deny all traffic from 0.0.0.0/0. Create an IAM policy for each user to prevent the use of any AWS Region other than ap-northeast-3.

E. Use AWS Config to activate managed rules to detect and alert for internet gateways and to detect and alert for new resources deployed outside of ap-northeast-3.

Correct Answer: AC

Community vote distribution



Comments

Six_Fingered_Jose **Highly Voted** 2 years, 7 months ago

Selected Answer: AC

agree with A and C

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps_examples_vpc.html#example_vpc_2

upvoted 20 times

cookieMr **Highly Voted** 1 year, 11 months ago

Selected Answer: AC

A. By using Control Tower, the company can enforce data residency guardrails and restrict internet access for VPCs and denies access to all Regions except the required ap-northeast-3 Region.

C. With Organizations, the company can configure SCPs to prevent VPCs from gaining internet access. By denying access to all

When Organizations, the company can configure rules to prevent VPCs from gaining internet access by denying access to all Regions except ap-northeast-3, the company ensures that VPCs can only be deployed in the specified Region.

Option B is incorrect because using rules in AWS WAF alone does not address the requirement of denying access to all AWS Regions except ap-northeast-3.

Option D is incorrect because configuring outbound rules in network ACLs and IAM policies for users can help restrict traffic and access, but it does not enforce the company's requirement of denying access to all Regions except ap-northeast-3.

Option E is incorrect because using AWS Config and managed rules can help detect and alert for specific resources and configurations, but it does not directly enforce the restriction of internet access or deny access to specific Regions.

upvoted 19 times

iamroyalty_k Most Recent 4 months ago

Selected Answer: CE

□ A. Control Tower does not directly prevent internet access; it only provides guardrails, but those can be bypassed in some cases.

□ C. SCPs (Service Control Policies) in AWS Organizations provide hard restrictions at the account level, making them more enforceable than Control Tower guardrails alone.

□ E. AWS Config ensures continuous monitoring and alerts, which Control Tower does not provide as effectively.

This ensures strong security controls, compliance enforcement, and real-time monitoring while maintaining AWS best practices.

upvoted 2 times

Skyskilo 4 months, 4 weeks ago

Selected Answer: CE

Option C provides centralized governance with SCPs, and Option E provides continuous monitoring and alerting for compliance. Together, these solutions meet the requirements of restricting internet access and ensuring usage of only the ap-northeast-3 region.

upvoted 1 times

PaulGa 8 months ago

Selected Answer: AC

Ans A, C - Control Tower with Organisations configured. The two go together

upvoted 2 times

ChymKuBoy 11 months, 4 weeks ago

Selected Answer: AC

AC for sure

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: AC

B: Irrelevant WAF

D: This is confusing so I'll ignore it.

E: Wrong product

A: Control Tower can have residency guard rails and block internet access.

C: SCP is like a duplicate of A IMHO but it stops admins from circumventing A as Org policies cannot be overridden by admins unless they are org admins.

Too moany assumptions

upvoted 3 times

BrijMohan08 1 year, 9 months ago

Selected Answer: AC

A. Use AWS Control Tower to implement data residency guardrails to deny internet access and deny access to all AWS Regions except ap-northeast-3.

C. Use AWS Organizations to configure service control policies (SCPs) that prevent VPCs from gaining internet access. Deny access to all AWS Regions except ap-northeast-3.

upvoted 3 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: AC

Use Control Tower to implement data residency guardrails and Service Control Policies (SCPS) to prevent VPCs from gaining internet access.

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: AC

AWS Control Tower guardrails and AWS Organizations SCPs provide centralized, automated mechanisms to enforce no internet connectivity for VPCs and restrict Region access to only ap-northeast-3.

upvoted 4 times

Abrar2022 2 years ago

Didn't know that SCPS (Service Control Policies) could be used to deny users internet access. Good to know. Always thought it's got controlling who can and can't access AWS Services.

upvoted 5 times

hicham0101 2 years, 1 month ago

Agree with A and C

<https://aws.amazon.com/blogs/aws/new-for-aws-control-tower-region-deny-and-guardrails-to-help-you-meet-data-residency-requirements/>

upvoted 2 times

yallahool 2 years, 2 months ago

I choose C and D.

For control tower, it can't be A because ap-northeast-3 doesn't support it!

Also, in the case of E, it is detection and warning, so it is difficult to prevent internet connection (although the view is a little obscure).

upvoted 1 times

micellemeloc 2 years, 1 month ago

I just check, now it's supported!!!

upvoted 3 times

notacert 2 years, 2 months ago

Selected Answer: AC

A and C

upvoted 1 times

datz 2 years, 2 months ago

Selected Answer: CD

C/D

A - CANNOT BE!!! AWS Control Tower is not available in ap-northeast-3! Check your

B - for sure no

C - SCPS (Service Control Policies)- For sure

D - Deny outbound rule to be place in prod and also IAM Policy to deny Users creating services in AP-Northeast3

E - it creates an alert, which means it happens but an alert is triggered. so I think it's not good either.

upvoted 2 times

jerry1 2 months, 3 weeks ago

I dont think region availability shall be a consideration in exam, which i think it would be crazy if aws intend to test your "knowledge" against regional availability of different service...

upvoted 1 times

dam 2 years, 1 month ago

False, Control Tower is in Osaka NorthEast 3

<https://docs.aws.amazon.com/controltower/latest/userguide/region-how.html>

upvoted 3 times

Kaireny54 2 years, 2 months ago

Selected Answer: CD

Control tower isn't available in AP-northeast-3 (only available in ap-northeast1 and 2 : <https://www.aws-services.info/controltower.html>)
For answer E, it creates an alert, which means it happens but an alert is triggered. so i think it's not good either.
That's why i would go for C and D
upvoted 2 times

dam 2 years, 1 month ago

False, Control Tower is in Osaka NorthEast 3
<https://docs.aws.amazon.com/controltower/latest/userguide/region-how.html>
upvoted 2 times

dam 2 years, 1 month ago

same page you posted:
ap-northeast-3 Asia Pacific (Osaka) 2023-04-20 <https://aws.amazon.com/controltower>
upvoted 2 times

Bmarodi 2 years ago

It's available now on the same link u pasted in earlier: ap-northeast-3 Asia Pacific (Osaka) 2023-04-20.
upvoted 2 times

WhericanIstart 2 years, 2 months ago

Selected Answer: CE

AWS Control tower is not available in ap-northeast-3!

<https://www.aws-services.info/controltower.html>
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #152

Topic 1

A company uses a three-tier web application to provide training to new employees. The application is accessed for only 12 hours every day. The company is using an Amazon RDS for MySQL DB instance to store information and wants to minimize costs.


What should a solutions architect do to meet these requirements?

- A. Configure an IAM policy for AWS Systems Manager Session Manager. Create an IAM role for the policy. Update the trust relationship of the role. Set up automatic start and stop for the DB instance.
- B. Create an Amazon ElastiCache for Redis cache cluster that gives users the ability to access the data from the cache when the DB instance is stopped. Invalidate the cache after the DB instance is started.
- C. Launch an Amazon EC2 instance. Create an IAM role that grants access to Amazon RDS. Attach the role to the EC2 instance. Configure a cron job to start and stop the EC2 instance on the desired schedule.
- D. Create AWS Lambda functions to start and stop the DB instance. Create Amazon EventBridge (Amazon CloudWatch Events) scheduled rules to invoke the Lambda functions. Configure the Lambda functions as event targets for the rules.

Most Voted

Correct Answer: D

Community vote distribution



Comments

study_aws1 Highly Voted 2 years, 7 months ago

<https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-lambda/>

It is option D. Option A could have been applicable had it been AWS Systems Manager State Manager & not AWS Systems Manager Session Manager

upvoted 40 times

123jhl0 Highly Voted 2 years, 7 months ago

Selected Answer: A

A is true for sure. "Schedule Amazon RDS stop and start using AWS Systems Manager" Steps in the documentation:

1. Configure an AWS Identity and Access Management (IAM) policy for State Manager.
2. Create an IAM role for the new policy.
3. Update the trust relationship of the role so Systems Manager can use it

3. Update the trust relationship so the role so Systems Manager can use it.

4. Set up the automatic stop with State Manager.

5. Set up the automatic start with State Manager.

<https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-systems-manager/>

upvoted 9 times

MatAlves 9 months, 2 weeks ago

The link you sent talks about "AWS Systems Manager State Manager", not "AWS Systems Manager Session Manager".

upvoted 2 times

NSA_Poker 1 year ago

(A) & (D) both look correct & both are supported by a 2-part article. First part describes (D), the second (A).

I don't have an issue with (A): Session vs State; I understand that State Manager is a capability of Systems Manager.

Because Lambda is not free beyond the free-tier limit & (A) is more complete with IAM notes unlike (D), I would choose (A) on the exam.

Schedule Amazon RDS stop and start using AWS Systems Manager (part-2)

<https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-systems-manager/>

Schedule Amazon RDS stop and start using AWS Lambda (part-1)

<https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-lambda/>

upvoted 2 times

Kien048 2 years, 7 months ago

And ofcause, D is working, so if A also right, the question is wrong.

upvoted 5 times

Kien048 2 years, 7 months ago

Look like State manager and Session manager use for difference purpose even both in same dashboard console.

upvoted 2 times

satyaamm Most Recent 4 months, 1 week ago

Selected Answer: D

If A had AWS System Manager State Manager then it would be A. Else Lambda is suitable here.

upvoted 1 times

lofzee 1 year ago

Selected Answer: D

D.

You need to use AWS Systems Manager State Manager, not Systems Manager Session Manager.

upvoted 2 times

ManikRoy 1 year, 1 month ago

Selected Answer: D

Agree with study_aws1 comment.

Lambda and event bridge solution is the correct answer.

The option A was only possible if it had mentioned System manager state manager

upvoted 2 times

phonglktv 1 year, 1 month ago

Selected Answer: C

we still pay for RDS even when the instance stops. So for cost-optimize -> C

upvoted 2 times

NSA_Poker 1 year ago

While the instance is stopped, you're charged for storage and backups, but not for the DB instance hours.

<https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-lambda/#:~:text=While%20the%20instance%20is%20stopped%2C%20you%E2%80%99re%20charged%20for%20storage%20a>

nd%20backups%2C%20but%20not%20for%20the%20DB%20instance%20hours

https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-lambda/
upvoted 2 times

LP0905 1 year, 3 months ago

Selected Answer: D

Although both A and D is a workable solution, the requirements is to minimum cost.

The benefits of automating the startup and shutdown of RDS DB instances using Lambda allows organizations to further reduce compute costs and simplify the administration of database environments that don't need to be running continuously.

<https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-lambda/>

For using system manager to accomplish the task works however keep in mind that although we're stopping the databases, the storage costs for the databases still apply.

<https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-systems-manager/>

Initially I also thought that A would be the correct answer however looking at the administration and cost I would go for D as a better solution instead.

upvoted 3 times

thewalker 1 year, 4 months ago

Selected Answer: A

To automatically shutdown an RDS instance during 09:00 PM to 09:00 AM and have it available between 09:00 AM to 09:00 PM, you can use AWS Systems Manager Maintenance Windows.

Create two Maintenance Windows - one to stop the RDS instance at 09:00 PM and another to start it at 09:00 AM.

For each Maintenance Window, select the "AWS-StopRDSInstance" and "AWS-StartRDSInstance" runbooks respectively and specify the cron expression for the schedule.

Tag the RDS instance with a name so it can be identified by the runbooks.

The runbooks will then automatically stop and start the RDS instance on the specified schedule without needing any manual intervention.

upvoted 1 times

thewalker 1 year, 4 months ago

This allows cost savings by shutting down the RDS instance during non-business hours while keeping it available during the day as per your requirements. Refer to the AWS documentation for more details on configuring Maintenance Windows and runbooks.

<https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/automatically-stop-and-start-an-amazon-rds-db-instance-using-aws-systems-manager-maintenance-windows.html>

https://repost.aws/questions/QUcVR5js8LSbOS_LE889KdIg/automatically-stop-and-start-an-amazon-rds-db-instance-in-a-cdk-app

upvoted 1 times

theochan 1 year, 5 months ago

Guys, we still have to pay for RDS instance even we stopped it, isn't it?

upvoted 3 times

NSA_Poker 1 year ago

While the instance is stopped, you're charged for storage and backups, but not for the DB instance hours.

<https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-lambda/#:~:text=While%20the%20instance%20is%20stopped%2C%20you%E2%80%99re%20charged%20for%20storage%20and%20backups%2C%20but%20not%20for%20the%20DB%20instance%20hours.>

upvoted 2 times

Ruffyt 1 year, 6 months ago

AWS Lambda functions can be used to start and stop RDS instances programmatically. EventBridge scheduled rules can trigger the Lambda functions at specified times daily. This allows fully automating the starting and stopping of RDS on a schedule to match usage patterns.

RDS billing is per hour when instance is running, so stopping when not in use significantly reduces costs. Using Lambda and EventBridge is simpler and more robust than cron jobs on EC2. ElastiCache and Systems Manager Session Manager are useful tools but do not directly address scheduled RDS start/stop.

upvoted 3 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: D

You can use AWS Lambda and Amazon EventBridge to schedule a Lambda function to stop and start the idle databases with specific tags to save on compute costs.

<https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-lambda/#:~:text=you%20to%20schedule%20a-,Lambda%20function,-to%20stop%20and%20start>

upvoted 3 times

lemur88 1 year, 9 months ago

Selected Answer: D

Here is the recommended solutions which describes choice D - <https://aws.amazon.com/blogs/database/save-costs-by-automating-the-start-and-stop-of-amazon-rds-instances-with-aws-lambda-and-amazon-eventbridge/>

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: D

AWS Lambda functions can be used to start and stop RDS instances programmatically. EventBridge scheduled rules can trigger the Lambda functions at specified times daily. This allows fully automating the starting and stopping of RDS on a schedule to match usage patterns. RDS billing is per hour when instance is running, so stopping when not in use significantly reduces costs. Using Lambda and EventBridge is simpler and more robust than cron jobs on EC2. ElastiCache and Systems Manager Session Manager are useful tools but do not directly address scheduled RDS start/stop.

upvoted 4 times

cookieMr 1 year, 11 months ago

Selected Answer: D

By using AWS Lambda functions triggered by Amazon EventBridge scheduled rules, the company can automate the start and stop actions for the Amazon RDS for MySQL DB instance based on the 12-hour access period. This allows them to minimize costs by only running the DB instance when it is needed.

Option A is not the most suitable solution because it refers to IAM policies for AWS Systems Manager Session Manager, which is primarily used for interactive shell access to EC2 instances and does not directly address the requirement of starting and stopping the DB instance.

Option B is not the most suitable solution because it suggests using Amazon ElastiCache for Redis as a cache cluster, which may not provide the desired cost optimization for the DB instance.

Option C is not the most suitable solution because launching an EC2 instance and configuring cron jobs to start and stop it does not directly address the requirement of minimizing costs for the Amazon RDS DB instance.

upvoted 4 times

Siva007 2 years ago

Selected Answer: D

I got this question in real exam!

upvoted 5 times

srijrao 1 year, 11 months ago

why we need more than one lambda function to start and stop DB instance? btw how many questions came from this site?

upvoted 2 times

ccmc 2 years, 1 month ago

State Manager, a capability of AWS Systems Manager

upvoted 2 times

Ankit_EC_ran 2 years, 1 month ago

Selected Answer: D

Option D is correct
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #153

Topic 1

A company sells ringtones created from clips of popular songs. The files containing the ringtones are stored in Amazon S3 Standard and are at least 128 KB in size. The company has millions of files, but downloads are infrequent for ringtones older than 90 days. The company needs to save money on storage while keeping the most accessed files readily available for its users.

Which action should the company take to meet these requirements MOST cost-effectively?

- A. Configure S3 Standard-Infrequent Access (S3 Standard-IA) storage for the initial storage tier of the objects.
- B. Move the files to S3 Intelligent-Tiering and configure it to move objects to a less expensive storage tier after 90 days.
- C. Configure S3 inventory to manage objects and move them to S3 Standard-Infrequent Access (S3 Standard-1A) after 90 days.
- D. Implement an S3 Lifecycle policy that moves the objects from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-1A) after 90 days. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****rjam** **Highly Voted** 2 years, 7 months ago**Selected Answer: D**

Answer D

Why Optoin D ?

The Question talks about downloads are infrequent older than 90 days which means files less than 90 days are accessed frequently. Standard-Infrequent Access (S3 Standard-IA) needs a minimum 30 days if accessed before, it costs more. So to access the files frequently you need a S3 Standard . After 90 days you can move it to Standard-Infrequent Access (S3 Standard-IA) as its going to be less frequently accessed

upvoted 40 times

MutiverseAgent 1 year, 10 months ago

I do not agree. The MOST cheaper option is B, because by choosing:

D) Files older than 90 days will live eternally in the S3 Infrequently access layer at \$0.0125/GB.

B) Using Intelligent-Tiering files older than 90 days can be moved DIRECTLY to the "Archive access tier" (Glacier instant retrieval) at \$0.004/GB, avoiding/skipping the "S3 Infrequently access layer". The question also seems to be according this

assumption as says "and configure it to move objects to a less expensive storage tier after 90 days".

<https://aws.amazon.com/s3/pricing/?nc=sn&loc=4>

upvoted 4 times

MutiverseAgent 1 year, 10 months ago

I am taking back my answer, the right is D) as the "Archive access tier" check present in the "Intelligent-Tiering Archive configurations" is for "S3 Glacier flexible retrieval" which is not instant retrieval.

upvoted 8 times

Itsmetanmay 1 year, 1 month ago

Archive Instant Access tier (automatic)

If an object is not accessed for 90 consecutive days, the object moves to the Archive Instant Access tier. The Archive Instant Access tier provides low latency and high-throughput performance. --- This says it moves to instant retrieval class. Price for this is lesser than standard IA. Answer should be B

upvoted 2 times

zeronine75 Highly Voted 2 years, 6 months ago

Selected Answer: B

B/D seems possible answer. But, I'll go with "B".

In the following table, S3 Intelligent-Tiering seems not so expansive than S3 Standard.

https://aws.amazon.com/s3/pricing/?nc1=h_ls

And, in the question "128KB" size is talking about S3 Intelligent-Tiering stuff.

upvoted 16 times

ruqui 2 years ago

have you tried to implement B? how do you configure Intelligent Tiering to move objects to a less expensive storage tier after 90 days? and which storage tier is this 'less expensive' ? the answer is clearly wrong ... correct answer is D

upvoted 5 times

FNJ1111 2 years, 5 months ago

also, there are probably several ringtones which aren't popular/used. Why keep them in S3 standard? The company would save money if s3 intelligent-tiering moves the unpopular ringtones to a more cost-effective tier than s3 standard.

upvoted 2 times

Wilson_S 2 years, 6 months ago

This link also has me going with "B." Specifying 128 KB in size is not a coincidence. <https://aws.amazon.com/s3/storage-classes/intelligent-tiering/>

upvoted 6 times

javitech83 2 years, 6 months ago

because of tha link it is D.

There are no retrieval charges in S3 Intelligent-Tiering. S3 Intelligent-Tiering has no minimum eligible object size, but objects smaller than 128 KB are not eligible for auto tiering. These smaller objects may be stored, but they'll always be charged at the Frequent Access tier

upvoted 1 times

javitech83 2 years, 6 months ago

oh sorry it states objects are bigger than 128 KB. B is correct

upvoted 2 times

Wajif 2 years, 6 months ago

S3 Intelligent tiering is used when the access frequency is not known. I think 128KB is a deflector.

upvoted 7 times

c12ab95 Most Recent 3 weeks, 4 days ago

Selected Answer: B

Solution is B.

Key Analysis

Infrequent Access After 90 Days: S3 Intelligent-Tiering automatically transitions objects not accessed for 90 days to the Archive Instant Access tier (\$0.004/GB-month), which is 68% cheaper than S3 Standard-IA (\$0.0125/GB-month). This aligns perfectly with the requirement to save costs for older ringtones.

No Retrieval Fees: Unlike S3 Standard-IA (which charges retrieval fees), Intelligent-Tiering has no retrieval costs for data in Frequent, Infrequent, or Archive Instant Access tiers .

Automated Tiering: Intelligent-Tiering eliminates manual lifecycle policies. Objects older than 90 days with no access are moved automatically, reducing operational overhead . For new or frequently accessed ringtones, objects remain in the Frequent Access tier (same cost as S3 Standard), ensuring immediate availability .

S3 Intelligent-Tiering provides deeper cost savings than Standard-IA for infrequently accessed data, automates tier transitions, and avoids retrieval fees, making it the optimal choice.

upvoted 1 times

Skyskilo 4 months, 4 weeks ago

Selected Answer: D

Option D is the most cost-effective solution as it leverages an S3 Lifecycle policy to automatically transition infrequently accessed files to S3 Standard-IA after 90 days, saving costs while meeting the company's requirements.

upvoted 1 times

Rcosmos 4 months, 4 weeks ago

Selected Answer: D

B. Mova os arquivos para o S3 Intelligent-Tiering:

O S3 Intelligent-Tiering é uma solução flexível, mas tem uma pequena taxa mensal para monitorar os objetos. No caso descrito, uma política de ciclo de vida seria mais econômica para o comportamento de acesso conhecido (raros acessos após 90 dias).

upvoted 1 times

[Removed] 6 months, 2 weeks ago

Selected Answer: D

Why D over the other options?

A. Configure S3 Standard-IA for initial storage: Not cost-effective for frequently accessed files during the first 90 days, as S3 Standard-IA is optimized for infrequent access.

B. Use S3 Intelligent-Tiering: While Intelligent-Tiering is a good choice for unpredictable access patterns, it incurs a monthly monitoring fee per object. Since the company's pattern is predictable (infrequent access after 90 days), this additional cost is unnecessary.

C. Use S3 inventory to manage objects manually: S3 inventory only generates reports and does not automate object transitions. Manually moving objects is less efficient and error-prone compared to automated lifecycle policies.

upvoted 3 times

PaulGa 8 months ago

Selected Answer: D

Ans D - as well explained by rjam (1 year, 11 months ago)

upvoted 2 times

Duckydoo 11 months, 3 weeks ago

Selected Answer: D

I am going with D instead of B because Intelligent-Tiering may work against you in this case. The company knows that the files are rarely accessed after 90 days. So even if a file is accessed on the 89th day, they would still want to "archive" that file the next day (again because they fully know the typical access behavior). However, Intelligent Tiering will reset the 90-day inactive clock starting on the 89th day (i.e. it will not archive the file until the 189th day assuming no other access).

upvoted 2 times

ChymKuBoy 11 months, 4 weeks ago

Selected Answer: D

D for sure

upvoted 1 times

lofzee 1 year ago

Selected Answer: D

Selected Answer: D

You can't configure Intelligent Tiering, that is the whole point of it. It's intelligent and moves things on its own so there is no 90 day configuration you can set.

Therefore D is the answer.

upvoted 1 times

lofzee 1 year ago

Actually scrap that. I'm changing my answer to B.

You can setup Intelligent Tiering to move data to an archive after 90 days.

https://docs.aws.amazon.com/AmazonS3/latest/userguide/storage-class-intro.html?icmpid=docs_amazons3_console#sc-dynamic-data-access

<https://aws.amazon.com/about-aws/whats-new/2021/11/s3-intelligent-tiering-archive-instant-access-tier/>

Also the 128KB size is perfect for Intelligent Tiering.

Tricky question indeed, one of those pointless ones that doesn't really make you a better or worse AWS Engineer

upvoted 1 times

NSA_Poker 1 year ago

Selected Answer: B

Some songs make comebacks. Lifecycle policy isn't intelligent enough to deal with a resurgence in popularity of a ringtone. 128KB in size makes files eligible for Intelligent-Tiering.

upvoted 1 times

ManikRoy 1 year, 1 month ago

Selected Answer: D

Clearly option D is the correct. This is not a trick question.

upvoted 1 times

MomenAWS 1 year, 1 month ago

Selected Answer: B

I think this is a kind of tricky question (B VS D) However:

The Q stated "most accessed files readily available for its users."

From this and correct me if I am wrong, I think B is a better solution since Amazon S3 Intelligent-Tiering working on "moving data to the most cost-effective access tier based on access frequency" while the S3 Standard-IA doesn't.

Ref: <https://aws.amazon.com/s3/storage-classes/>

upvoted 1 times

richiexamaws 1 year, 2 months ago

Selected Answer: B

B. Move the files to S3 Intelligent-Tiering and configure it to move objects to a less expensive storage tier after 90 days.

Explanation:

S3 Intelligent-Tiering is a storage class that automatically moves objects between two access tiers: frequent access and infrequent access, based on access patterns.

By configuring S3 Intelligent-Tiering to move objects to a less expensive storage tier after 90 days of infrequent access, the company can save money on storage costs while ensuring that the most accessed files remain readily available.

This approach is more cost-effective than using S3 Standard-Infrequent Access (S3 Standard-IA) because it automatically adjusts storage tiers based on access patterns without the need for manual configuration or management.

upvoted 1 times

Kanagarajd 1 year, 3 months ago

Selected Answer: D

Right answer is D

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: D

D is cheapest and managed by S3 Lifecycle policy

A: Not readily available

C: Wrong product

B: No choice of '90 days' so you'll be paying for Intelligent Tiering unnecessarily for files to drop out of frequent access after the first 90 days.

upvoted 1 times

Firdous586 1 year, 5 months ago

B is the correct answer Kindly follow the below link for more information as proof

<https://aws.amazon.com/s3/storage-classes/>

upvoted 1 times

awsgeek75 1 year, 4 months ago

So D is the correct answer as IA is cheaper than Intelligent tier.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #154

Topic 1


A company needs to save the results from a medical trial to an Amazon S3 repository. The repository must allow a few scientists to add new files and must restrict all other users to read-only access. No users can have the ability to modify or delete any files in the repository. The company must keep every file in the repository for a minimum of 1 year after its creation date.

Which solution will meet these requirements?

- A. Use S3 Object Lock in governance mode with a legal hold of 1 year.
- B. Use S3 Object Lock in compliance mode with a retention period of 365 days. **Most Voted**
- C. Use an IAM role to restrict all users from deleting or changing objects in the S3 bucket. Use an S3 bucket policy to only allow the IAM role.
- D. Configure the S3 bucket to invoke an AWS Lambda function every time an object is added. Configure the function to track the hash of the saved object so that modified objects can be marked accordingly.

Correct Answer: B

Community vote distribution



Comments

elmogy **Highly Voted** 2 years ago

Selected Answer: B

B,
The key is "No users can have the ability to modify or delete any files" and compliance mode supports that.
I remember it this way: (governance is like government, they set the rules but they can allow some people to break it :D)
upvoted 60 times

Burrito69 1 year, 2 months ago

I liked that thought of yours.. can you do more of these please? Thank you
upvoted 5 times

Praewwara 1 year, 5 months ago

Amazon S3 Object Lock

1. Governance mode - Only users with special permissions can overwrite, delete, or alter object lock settings

2. Compliance mode - No user, including the root user in an AWS account, can overwrite, delete, or alter object lock settings
upvoted 12 times

Qjb8m9h Highly Voted 2 years, 7 months ago

Answer : B

Reason: Compliance Mode. The key difference between Compliance Mode and Governance Mode is that there are NO users that can override the retention periods set or delete an object, and that also includes your AWS root account which has the highest privileges.

upvoted 22 times

Zerotn3 2 years, 5 months ago

How about: The repository must allow a few scientists to add new files

upvoted 1 times

JayBee65 2 years, 5 months ago

Adding is not the same as changing :)

upvoted 8 times

abhishek2021 2 years ago

Compliance mode controls the object life span after creation.
how this option restricts all scientists from adding new file? please explain.

upvoted 4 times

satyaamm Most Recent 4 months, 1 week ago

Selected Answer: B

Compliance mode supports the feature of "No users can delete or modify the files".

upvoted 1 times

PaulGa 8 months ago

Selected Answer: B

Ans B - Compliance mode... but not sure that answers "allow only a few scientists to add new files" ...?

upvoted 3 times

toyaji 10 months, 1 week ago

Selected Answer: B

First of all, Regal hold has no expiration before you remove it. So A makes no sense.

After that Governance mode is breakable with permission, but Compliance mode is not even for root user cannot delete it.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.html>

upvoted 2 times

ChymKuBoy 11 months, 3 weeks ago

Selected Answer: B

B for sure

upvoted 1 times

ManikRoy 1 year, 1 month ago

Selected Answer: B

I almost chose A for this deciving line lol but it would be compliance mode as no user should be able to change objects:-

The repository must allow a few scientists to add new files and must restrict all other users to read-only access.

upvoted 2 times

demigodnyi 1 year, 4 months ago

Can someone please explain why the answer is not A. It said that The repository must allow a few scientists to add new files. So, i think some user must have permission to change it.

upvoted 1 times

NSA_Poker 1 year ago

Legal hold – A legal hold provides the same protection as a retention period, but it HAS NO EXPIRATION DATE. Instead, a legal hold remains in place until you explicitly remove it. Legal holds are independent from retention periods and are PLACED ON INDIVIDUAL OBJECT VERSIONS.

upvoted 2 times

ManikRoy 1 year, 1 month ago

Write Once Read Many - in compliance mode there is no restriction in adding but only changing existing objects

upvoted 2 times

pentium75 1 year, 5 months ago

Unsure, B would meet the "must keep every file for a minimum of 1 year" requirement. (In theory C would too if you ignore the root user, but administrators could remove the policy.) But what about the 'a few scientists must be able to add new files'? None of the options mentions permissions for a special group.

upvoted 3 times

awsgeek75 1 year, 4 months ago

agree that something is missing for "some users".

ACD are not going to work flat out so B looks like right answer but with some language issues either in the question or the answer.

upvoted 1 times

LoXoL 1 year, 5 months ago

Agree. It looks like it's missing sth here.

upvoted 1 times

Ruffyt 1 year, 6 months ago

Both Compliance & Governance mode protect objects against being deleted or changed. But in Governance mode some people can have special permissions. In this question, no user can delete or modify files; so the answer is Compliance mode only. Neither of these modes restrict user from adding new files.

upvoted 3 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: B

Compliance Mode best suits this scenario because once an object is locked in compliance mode, its retention mode can't be changed, and its retention period can't be shortened.

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

B) seems to be the right option, because: Both option A) & B) allow to:

- Scientists add new files & other users read-only access.
- Keep files for a minimum of 1 year

Only option B allows to:

- Disable all users the ability to modify or delete any file.

If A) were the correct option some scientis will be able to modify files, as if they were in charge of put an object lock same permission would allow them to remove the lock and consequently delete the file.

upvoted 2 times

MutiverseAgent 1 year, 10 months ago

Selected Answer: B

B) seems to be the right option, because: Both option A) & B) allow to:

- Scientists add new files & other users read-only access.
- Keep files for a minimum of 1 year

Only option B allows to:

- Disable all users the ability to modify or delete any file.

If A) were the correct option some scientis will be able to modify files, as if they were in charge of put an object lock same permission would allow them to remove the lock and consequently delete the file.

upvoted 1 times

cookieMr 1 year, 11 months ago

Selected Answer: B

S3 Object Lock provides the necessary features to enforce immutability and retention of objects in an S3. Compliance mode ensures that the locked objects cannot be deleted or modified by any user, including those with write access. By setting a retention period of 365 days, the company can ensure that every file in the repository is kept for a minimum of 1 year after its creation date.

A does not provide the same level of protection as compliance mode. In governance mode, there is a possibility for authorized users to remove the legal hold, potentially allowing objects to be modified or deleted.

C can restrict users from deleting or changing objects, but it does not enforce the retention period requirement. It also does not provide the same level of immutability and protection against accidental or malicious modifications.

D does not address the requirement of preventing users from modifying or deleting files. It provides a mechanism for tracking changes but does not enforce the desired access restrictions or retention period.

upvoted 4 times

norris81 2 years ago

Am I the only one to worry about leap years ?

upvoted 1 times

cheese929 2 years, 1 month ago

Selected Answer: B

In compliance mode, a protected object version can't be overwritten or deleted by any user, including the root user in your AWS account. When an object is locked in compliance mode, its retention mode can't be changed, and its retention period can't be shortened. Compliance mode helps ensure that an object version can't be overwritten or deleted for the duration of the retention period.

In governance mode, users can't overwrite or delete an object version or alter its lock settings unless they have special permissions. With governance mode, you protect objects against being deleted by most users, but you can still grant some users permission to alter the retention settings or delete the object if necessary.

In Governance mode, Objects can be deleted by some users with special permissions, this is against the requirement.

upvoted 2 times

dam 2 years, 1 month ago

Selected Answer: B

its B, legal hold has no retention

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #155

Topic 1


A large media company hosts a web application on AWS. The company wants to start caching confidential media files so that users around the world will have reliable access to the files. The content is stored in Amazon S3 buckets. The company must deliver the content quickly, regardless of where the requests originate geographically.

Which solution will meet these requirements?

- A. Use AWS DataSync to connect the S3 buckets to the web application.
- B. Deploy AWS Global Accelerator to connect the S3 buckets to the web application.
- C. Deploy Amazon CloudFront to connect the S3 buckets to CloudFront edge servers. **Most Voted**
- D. Use Amazon Simple Queue Service (Amazon SQS) to connect the S3 buckets to the web application.

Correct Answer: C

Community vote distribution



Comments

rjam **Highly Voted** 2 years, 7 months ago

key :caching
Option C
upvoted 15 times

Guru4Cloud **Highly Voted** 1 year, 10 months ago

Selected Answer: C

The reasons are:

Amazon CloudFront is a content delivery network (CDN) that caches content at edge locations around the world. Connecting the S3 buckets containing the media files to CloudFront will cache the content at global edge locations. This provides fast reliable access to users everywhere by serving content from the nearest edge location. CloudFront integrates tightly with S3 for secure, durable storage. Global Accelerator improves availability and performance for TCP/UDP traffic, not HTTP-based content delivery. DataSync and SQS are not technologies for a global CDN like CloudFront.

upvoted 8 times

satyaamm **Most Recent** 4 months, 1 week ago

Selected Answer: C