

Buruguduystunstugudunstuy 2 years, 5 months ago

I would go with Option B since it meets the company's requirements and is the most suitable solution.

By creating an Amazon RDS MySQL DB instance with Multi-AZ functionality enabled, the solutions architect will ensure that data is automatically synchronously replicated across multiple AZs within the same Region. This provides high availability and data durability, minimizing the risk of data loss and ensuring that every transaction is stored on at least two nodes.

upvoted 3 times

JayBee65 2 years, 5 months ago

Option B is not incorrect: "The primary DB instance is synchronously replicated across Availability Zones to a standby replica to provide data redundancy and minimize latency spikes during system backups" from <https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZSingleStandby.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #183

Topic 1

A company is building a new dynamic ordering website. The company wants to minimize server maintenance and patching. The website must be highly available and must scale read and write capacity as quickly as possible to meet changes in user demand.

Which solution will meet these requirements?

A. Host static content in Amazon S3. Host dynamic content by using Amazon API Gateway and AWS Lambda. Use Amazon DynamoDB with on-demand capacity for the database. Configure Amazon CloudFront to deliver the website content.

Most Voted


B. Host static content in Amazon S3. Host dynamic content by using Amazon API Gateway and AWS Lambda. Use Amazon Aurora with Aurora Auto Scaling for the database. Configure Amazon CloudFront to deliver the website content.

C. Host all the website content on Amazon EC2 instances. Create an Auto Scaling group to scale the EC2 instances. Use an Application Load Balancer to distribute traffic. Use Amazon DynamoDB with provisioned write capacity for the database.

D. Host all the website content on Amazon EC2 instances. Create an Auto Scaling group to scale the EC2 instances. Use an Application Load Balancer to distribute traffic. Use Amazon Aurora with Aurora Auto Scaling for the database.

Correct Answer: A

Community vote distribution



Comments

romko Highly Voted 2 years, 6 months ago

Selected Answer: A

A - is correct, because Dynamodb on-demand scales write and read capacity

B - Aurora auto scaling scales only read replicas

upvoted 52 times

klayytech 2 years, 2 months ago

That's not correct. Amazon Aurora with Aurora Auto Scaling can scale both read and write replicas. Is there anything else you would like me to help you with?

upvoted 7 times

Duckydoo 12 months ago

Could you point us to a source where it says that Aurora Auto Scaling can scale write replicas? The AWS documentation specifically mentions that it supports only read replicas (e.g. <https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.Integrating.AutoScaling.html>):

To meet your connectivity and workload requirements, Aurora Auto Scaling dynamically adjusts the number of Aurora Replicas (reader DB instances) provisioned for an Aurora DB cluster. Aurora Auto Scaling is available for both Aurora MySQL and Aurora PostgreSQL.

upvoted 2 times

Yadav_Sanjay 1 year, 11 months ago

That's why Dynamo DB is best suited option

upvoted 3 times

SteveNguyen 6 months ago

Dynamo is serverless, so you don't need to worry about maintenance and patching.

upvoted 1 times

Yadav_Sanjay 1 year, 11 months ago

Correct...Both can serve purpose but note the keyword "must scale read and write capacity as quickly as possible to meet changes in user demand". DynamoDB can scale quickly than Aurora. Remember "PUSH BUTTON SCALING FEATURE" of Dynamo DB.

upvoted 7 times

Manlikeleke **Highly Voted** 2 years, 6 months ago

please is this dump enough to pass the exam?

upvoted 15 times

Bobbybash 2 years, 6 months ago

I HOPE SO

upvoted 9 times

LuckyAro 2 years, 4 months ago

You can tell us now ? Going by the date of your post I guess you would have challenged the exam by now ? so how did it go ?

upvoted 9 times

soufiyane 1 year, 2 months ago

did you pass ?

upvoted 2 times

tch **Most Recent** 3 months ago

Selected Answer: A

Aurora is more complex to set up and manage....

DynamoDB is simpler to use, less configuration needed, the most important is highly scalable, auto-scaling

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/AutoScaling.html>

upvoted 1 times

satyaamm 4 months ago

Selected Answer: A

DynamoDB provides the best read and write capacity here.

upvoted 1 times

Dharmarajan 4 months, 1 week ago

Selected Answer: A

I also initially thought Aurora is better suited since it gives auto scaling as well as auto scaling. However, it doesn't auto scale for writer instances.

there is a note in Aurora doc:

"Aurora Auto Scaling doesn't apply to the workload on the writer DB instance. Aurora Auto Scaling helps with the workload only on the reader instances."
<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.Integrating.AutoScaling.html>

Hence the right answer is A. But man, how devious! a small detail can derail you from making the right choice!
upvoted 3 times

devarajchidambaram 5 months ago

Selected Answer: A

Dynamodb - No need to provision, patch, or manage servers.
upvoted 1 times

ManikRoy 1 year, 1 month ago

Selected Answer: A

Dynamo DB Push Scaling
upvoted 2 times

Cloud_A 1 year, 5 months ago

Selected Answer: A

[https://aws.amazon.com/blogs/database/how-to-determine-if-amazon-dynamodb-is-appropriate-for-your-needs-and-then-plan-your-migration/#:~:text=Are%20working%20with%20an%20online%20transaction%20processing%20\(OLTP\)%20workload.%20High%2Dperformance%20reads%20and%20writes%20are%20easy%20to%20manage%20with%20DynamoDB%2C%20and%20you%20can%20expect%20performance%20that%20is%20effectively%20constant%20across%20widely%20varying%20loads.](https://aws.amazon.com/blogs/database/how-to-determine-if-amazon-dynamodb-is-appropriate-for-your-needs-and-then-plan-your-migration/#:~:text=Are%20working%20with%20an%20online%20transaction%20processing%20(OLTP)%20workload.%20High%2Dperformance%20reads%20and%20writes%20are%20easy%20to%20manage%20with%20DynamoDB%2C%20and%20you%20can%20expect%20performance%20that%20is%20effectively%20constant%20across%20widely%20varying%20loads.)
upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

C,D are out due to EC2 scaling which is not ideal for static content scaling.
A and B are logical choices. B uses Aurora which is more for relational database and comes with the baggage and limitations of RDBMS scaling. DynamoDB (no SQL) is easier to scale for both read and write. A is simply better than B for an ordering website so that is the better option. Note that B would have been good if A wasn't a choice.
upvoted 3 times

ChaBum 9 months, 2 weeks ago

Where do you see "static content scaling"?
In the question, I can read "dynamic ordering website", I fail to see how dynamic website can be hosted on S#.
upvoted 2 times

tom_cruise 1 year, 8 months ago

Selected Answer: A

dynamodb is serverless
upvoted 5 times

Angryasianxd 1 year, 9 months ago

Selected Answer: A

Hi all! The answer is A and NOT B on this one as the company is building an ordering website (OLTP). DynamoDB's high performance read and writes are perfect for an OLTP use case.

<https://aws.amazon.com/blogs/database/how-to-determine-if-amazon-dynamodb-is-appropriate-for-your-needs-and-then-plan-your-migration/>
upvoted 4 times

n0pz 1 year, 9 months ago

S3 is discarded since the question says: A company is building a new dynamic ordering website,
upvoted 2 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: A

minimize server maintenance and patching, highly available, scale read and write = serverless = Amazon S3, Amazon API Gateway, AWS Lambda, Amazon DynamoDB

upvoted 3 times

DebAwsAccount 1 year, 9 months ago

Selected Answer: A

Key phrase in the Question is must scale read and write capacity. Aurora is only for Read.

Amazon DynamoDB has two read/write capacity modes for processing reads and writes on your tables:

On-demand

Provisioned (default, free-tier eligible)

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ReadWriteCapacityMode.html>

upvoted 4 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: A

Minimize maintenance & Patching = Serverless

S3, DynamoDB are serverless

upvoted 2 times

ravindrabagale 1 year, 10 months ago

Minimize maintenance & Patching = Serverless services

Serverless services with no sql database is perfect combination

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: A

B. This solution leverages serverless technologies like API Gateway and Lambda for hosting dynamic content, reducing server maintenance and patching. Aurora with Aurora Auto Scaling provides a highly available and scalable database solution. Hosting static content in S3 and configuring CloudFront for content delivery ensures high availability and efficient scaling.

A. Using DynamoDB with on-demand capacity may provide scalability, but it does not offer the same level of flexibility and performance as Aurora. Additionally, it does not address the hosting of dynamic content using serverless technologies.

C. Hosting all the website content on EC2 instances requires server maintenance and patching. While using ASG and an ALB helps with availability and scalability, it does not minimize server maintenance as requested.

D. Hosting all the website content on EC2 instances introduces server maintenance and patching. Using Aurora with Aurora Auto Scaling is a good choice for the database, but it does not address the need to minimize server maintenance and patching for the overall infrastructure.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #184

Topic 1

A company has an AWS account used for software engineering. The AWS account has access to the company's on-premises data center through a pair of AWS Direct Connect connections. All non-VPC traffic routes to the virtual private gateway.


A development team recently created an AWS Lambda function through the console. The development team needs to allow the function to access a database that runs in a private subnet in the company's data center.

Which solution will meet these requirements?

- A. Configure the Lambda function to run in the VPC with the appropriate security group. **Most Voted**
- B. Set up a VPN connection from AWS to the data center. Route the traffic from the Lambda function through the VPN.
- C. Update the route tables in the VPC to allow the Lambda function to access the on-premises data center through Direct Connect.
- D. Create an Elastic IP address. Configure the Lambda function to send traffic through the Elastic IP address without an elastic network interface.

Correct Answer: A

Community vote distribution



Comments

Gil80 **Highly Voted** 2 years, 6 months ago

Selected Answer: A

To configure a VPC for an existing function:

1. Open the Functions page of the Lambda console.
2. Choose a function.
3. Choose Configuration and then choose VPC.
4. Under VPC, choose Edit.
5. Choose a VPC, subnets, and security groups. <-- **That's why I believe the answer is A**.

Note:

If your function needs internet access, use network address translation (NAT). Connecting a function to a public subnet doesn't give it internet access or a public IP address.

invited 24 times

upvoted 24 times

markw92 1 year, 12 months ago

The question says on-prem database...how do we create a SG for that instance in AWS? C make sense. my 2 cents..

upvoted 7 times

SSadiq 1 year ago

SG is for Lambda and not for the on-prem database. A is the correct option

upvoted 2 times

AZ_Master 1 year, 6 months ago

A is correct. To configure SG for Lambda , go to Lambda function -> Configure -> Edit VPC and scroll down to see "security groups" where you can configure Lambda for VPC.

Also see here

<https://repost.aws/questions/QUSaj1a6jBQ92Kp56klbZFNw/aws-lambda-to-on-premise-via-direct-connect-and-aws-privatelink>

upvoted 2 times

javitech83 Highly Voted 2 years, 6 months ago

Selected Answer: A

it is A. C is not correct at all as in the question it metions that the VPC already has connectivity with on-premises

upvoted 11 times

LuckyAro 2 years, 4 months ago

C says to "update the route table" not create a new connection. C is correct.

upvoted 5 times

ruqui 2 years ago

C is wrong. Lambda can't connect by default to resources in a private VPC, so you have to do some specific setup steps to run in a private VPC, Answer A is correct

upvoted 3 times

Adios_Amigo 2 years, 1 month ago

No need to do route updates. This is because the route to the destination on-premises is already set.

upvoted 4 times

jerryI 2 months, 3 weeks ago

the question didnt mention where the direct connect is set up but mention that all non vpc traffic are private so i think it makes sense you need to config route table

and i think in C, update the vpc already imply that the lambda is going to be run in that vpc

upvoted 1 times

ChhatwaniB Most Recent 2 months, 3 weeks ago

Selected Answer: C

Same queston in exam topic #524 and highly voted answer is C there

Since database on company's data center can not have security group. my vote goes to C

upvoted 1 times

Dharmarajan 4 months, 1 week ago

Selected Answer: A

The language is confusing, but the option A essentially says that "Configure the Lambda to execute in your own VPC (instead of AWS's own Lambda VPC) and since your VPC already is configured with the route to On prem data center, it will be able to access the on prem data center resources.

upvoted 1 times

FlyingHawk 4 months, 1 week ago

Selected Answer: A

By default, Lambda functions operate outside the customer's VPC. To access on-premises resources via Direct Connect, the Lambda function must be deployed inside the VPC.

Once attached to the VPC, Lambda can use the VPC's networking setup (route tables, security groups, and the virtual private gateway) to communicate with the on-premises database over Direct Connect.

upvoted 2 times

ensbrvsnss 9 months, 3 weeks ago

Selected Answer: C

either A or C

upvoted 1 times

jatric 11 months, 1 week ago

Selected Answer: C

C is correct as lambda already in VPC and AWS account already has connection setup with on-premise database in private subnet

upvoted 2 times

lofzee 1 year ago

Selected Answer: A

B,C,D dont have any logic behind them.

A is the most logical answer as you need to connect a function to a VPC. The VPC will be connected to the on-prem database.

upvoted 2 times

MehulKapadia 1 year, 2 months ago

Selected Answer: A

Answer A: During Lambda function creation select "Advanced Settings" select "Enable VPC", this will allow you to select VPC, Subnets and SecurityGroup for your Lambda function. This is the way Lambda can get controlled access to resouces in your VPC.

Default Lambda Settings:

When you create a Lambda function without specifying a VPC, the Lambda function does not get associated with any particular VPC. By default, Lambda functions are not deployed within a VPC and do not have access to resources within a VPC, such as EC2 instances, RDS databases, or ElastiCache clusters, unless you explicitly configure the Lambda function to connect to a VPC.

upvoted 5 times

Uzbekistan 1 year, 2 months ago

Selected Answer: C

Update the route tables in the VPC to allow the Lambda function to access the on-premises data center through Direct Connect.

By updating the route tables in the VPC to allow the Lambda function to access the on-premises data center through Direct Connect, is the most appropriate solution. By updating the route tables, you can specify the route for traffic from the Lambda function to the IP address range of the on-premises data center via the Direct Connect connection. This ensures that the Lambda function can securely communicate with the database in the private subnet of the data center.

upvoted 2 times

awsgeek75 1 year, 4 months ago

Every time I read this question the badly phrased options make no sense at all. I now want to vote for A but it makes no sense.

Question says: All non-VPC traffic routes to the virtual private gateway

So Lambda is technically a non VPC traffic too. This means it already goes through the VPGW but we don't know what it connects. Assuming it connect the data-centre to AWS then A makes sense. BUT all this is based on different interpretation now for me.

upvoted 9 times

pentium75 1 year, 5 months ago

Selected Answer: A

The wording is strange because technically, the Lambda function does not "run in the VPC", rather it is connected to the VPC, but otherwise A is what relevant documentation says - connect the Lambda function to the VPN and allow traffic in the security group.

Not B, we have Direct Connect, no need for VPN.

Not C, route is already in place. And route alone does not help - the "route tables in the VPC" are completely irrelevant as long

NOT C, route is already in place. And route alone does not help - the route tables in the VPC are completely irrelevant as long as we don't connect the Lambda function to the VPC.

Not D, an "Elastic IP address" is always connected to an "elastic network interface", such is created automatically with A.
upvoted 7 times

Kanagarajd 1 year, 3 months ago

I agree with explanation!
upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

The question and options are very badly worded so it makes C a possible candidate (unconvincingly though!).

B: VPN is not needed as Direct Connect is already there

D: Irrelevant

A is too generic (appropriate security group for what?) Lambda has fixed VPC or ENI

C is logically relevant

upvoted 1 times

pentium75 1 year, 5 months ago

A says "configure the Lambda function to RUN IN the VPC", but "a Lambda function ALWAYS runs inside a VPC owned by the Lambda service" (<https://docs.aws.amazon.com/lambda/latest/dg/foundation-networking.html>). "You can configure a Lambda function to CONNECT TO private subnets in a virtual private cloud (VPC) in your AWS account", but "connect to" is not the same as "run in" (<https://docs.aws.amazon.com/lambda/latest/dg/configuration-vpc.html>). Otherwise A would make sense (you CAN assign a security group to the Elastic Network Interface that Lambda uses to connect to your VPC).

upvoted 2 times

pentium75 1 year, 5 months ago

B We already have Direct Connect, so why set up VPN

C doesn't make sense because "all non-VPC traffic [already] routes to the virtual private gateway" (which is obviously connected to the Direct Connect gateway), so why should you "update the route tables"?

D sounds plausible; however, an Elastic IP address is associated with an Elastic Network Interface (though that is automatically provided by AWS). So the "without an elastic network interface" makes D wrong.

My best guess is that there's a typo or misunderstanding in the answers. It's either A but it should read "connect to the VPC" instead of "run in the VPC", or it's D but it should read "without CREATING an elastic network interface" or "WITH an elastic network interface".

upvoted 2 times

xdkonorek2 1 year, 7 months ago

Selected Answer: C

it's not A:

A Lambda function always runs inside a VPC owned by the Lambda service.

<https://docs.aws.amazon.com/lambda/latest/dg/foundation-networking.html>

upvoted 2 times

liux99 1 year, 7 months ago

The answer is C. The question is to allow lambda to access the database running in private subnet in the corporate data center. The only connectivity with the data center is Direct connect.

upvoted 2 times

lgogor 1 year, 8 months ago

Answer C is correct:

<https://repost.aws/questions/QU5aj1a6jBQ92Kp56klbZFNw/aws-lambda-to-on-premise-via-direct-connect-and-aws-privatelink>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #185

Topic 1

A company runs an application using Amazon ECS. The application creates resized versions of an original image and then makes Amazon S3 API calls to store the resized images in Amazon S3.

How can a solutions architect ensure that the application has permission to access Amazon S3?

- A. Update the S3 role in AWS IAM to allow read/write access from Amazon ECS, and then relaunch the container.
- B. Create an IAM role with S3 permissions, and then specify that role as the taskRoleArn in the task definition. **Most Voted**
- C. Create a security group that allows access from Amazon ECS to Amazon S3, and update the launch configuration used by the ECS cluster.
- D. Create an IAM user with S3 permissions, and then relaunch the Amazon EC2 instances for the ECS cluster while logged in as this account.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Buruguduystunstugudunstuy **Highly Voted** 1 year, 5 months ago

Selected Answer: B

To ensure that an Amazon Elastic Container Service (ECS) application has permission to access Amazon Simple Storage Service (S3), the correct solution is to create an AWS Identity and Access Management (IAM) role with the necessary S3 permissions and specify that role as the taskRoleArn in the task definition for the ECS application.

Option B, creating an IAM role with S3 permissions and specifying that role as the taskRoleArn in the task definition, is the correct solution to meet the requirement.

upvoted 14 times

Buruguduystunstugudunstuy 1 year, 5 months ago

Option A, updating the S3 role in IAM to allow read/write access from ECS and relaunching the container, is not the correct solution because the S3 role is not associated with the ECS application.

Option C, creating a security group that allows access from ECS to S3 and updating the launch configuration used by the ECS cluster, is not the correct solution because security groups are used to control inbound and outbound traffic to resources, and do not grant permissions to access resources.

do not grant permissions to access resources.

Option D, creating an IAM user with S3 permissions and relaunching the EC2 instances for the ECS cluster while logged in as this account, is not the correct solution because it is generally considered best practice to use IAM roles rather than IAM users to grant permissions to resources.

upvoted 9 times

Guru4Cloud Most Recent 9 months ago

Selected Answer: B

B. Create an IAM role with S3 permissions, and then specify that role as the taskRoleArn in the task definition

upvoted 4 times

cookieMr 11 months, 3 weeks ago

Selected Answer: B

Option B: Create an IAM role with S3 permissions and specify that role as the taskRoleArn in the task definition. This approach allows the ECS task to assume the specified role and gain the necessary permissions to access Amazon S3.

Option A is incorrect because updating the S3 role in IAM and relaunching the container does not associate the updated role with the ECS task.

Option C is incorrect because creating a security group that allows access from Amazon ECS to Amazon S3 does not grant the necessary permissions to the ECS task.

Option D is incorrect because creating an IAM user with S3 permissions and relaunching the EC2 instances for the ECS cluster does not associate the IAM user with the ECS task.

upvoted 4 times

dydzah 1 year ago

<https://repost.aws/knowledge-center/ecs-fargate-access-aws-services>

upvoted 2 times

k1kavi1 1 year, 5 months ago

Selected Answer: B

<https://www.examttopics.com/discussions/amazon/view/27954-exam-aws-certified-solutions-architect-associate-saa-c02/>

<https://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/aws-resource-ecs-taskdefinition.html>

upvoted 3 times

techhb 1 year, 5 months ago

Selected Answer: B

The short name or full Amazon Resource Name (ARN) of the AWS Identity and Access Management role that grants containers in the task permission to call AWS APIs on your behalf.

upvoted 3 times

BENICE 1 year, 5 months ago

Option B

upvoted 2 times

career360guru 1 year, 5 months ago

Selected Answer: B

Option B.

upvoted 3 times

k1kavi1 1 year, 6 months ago

Selected Answer: B

Agreed

upvoted 1 times

lighrz 1 year, 6 months ago

Selected Answer: B

B is the best answer
upvoted 1 times

Wpcorgan 1 year, 6 months ago

B is correct
upvoted 1 times

taer 1 year, 6 months ago

Selected Answer: B

The answer is B.
upvoted 1 times

Nigma 1 year, 6 months ago

B is the answer
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #186

Topic 1

A company has a Windows-based application that must be migrated to AWS. The application requires the use of a shared Windows file system attached to multiple Amazon EC2 Windows instances that are deployed across multiple Availability Zone:

What should a solutions architect do to meet this requirement?

- A. Configure AWS Storage Gateway in volume gateway mode. Mount the volume to each Windows instance.
- B. Configure Amazon FSx for Windows File Server. Mount the Amazon FSx file system to each Windows instance. **Most Voted**
- C. Configure a file system by using Amazon Elastic File System (Amazon EFS). Mount the EFS file system to each Windows instance.
- D. Configure an Amazon Elastic Block Store (Amazon EBS) volume with the required size. Attach each EC2 instance to the volume. Mount the file system within the volume to each Windows instance.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Nigma **Highly Voted** 1 year, 6 months ago

Correct is B
FSx --> shared Windows file system (SMB)
EFS --> Linux NFS
upvoted 10 times

satyaamm **Most Recent** 4 months ago

Selected Answer: B

FSx is most suitable here for windows file systems.
upvoted 1 times

TariqKipkemei 9 months ago

Selected Answer: B

Windows file system = Amazon FSx for Windows File Server
upvoted 4 times

Guru4Cloud 9 months, 4 weeks ago

Selected Answer: B

Configure Amazon FSx for Windows File Server. Mount the Amazon FSx file system to each Windows instance.
upvoted 2 times

cookieMr 11 months, 3 weeks ago

Selected Answer: B

Option B: Configure Amazon FSx for Windows File Server. This service provides a fully managed Windows file system that can be easily shared across multiple EC2 Windows instances. It offers high performance and supports Windows applications that require file storage.

Option A is incorrect because AWS Storage Gateway in volume gateway mode is not designed for shared file systems.

Option C is incorrect because while Amazon EFS can be mounted to multiple instances, it is a Linux-based file system and may not be suitable for Windows applications.

Option D is incorrect because attaching and mounting an Amazon EBS volume to multiple instances simultaneously is not supported.

upvoted 3 times

Bmarodi 1 year ago

Selected Answer: B

Option B is right answer.

upvoted 1 times

k1kavi1 1 year, 5 months ago

Selected Answer: B

References :

<https://www.examtips.com/discussions/amazon/view/28006-exam-aws-certified-solutions-architect-associate-saa-c02/>

<https://docs.aws.amazon.com/AmazonECS/latest/developerguide/wfsx-volumes.html>

upvoted 2 times

techhb 1 year, 5 months ago

Selected Answer: B

EFS is not compatible with Windows.

<https://pilotcoresystems.com/insights/ebs-efs-fsx-s3-how-these-storage-options-differ/#:~:text=EFS%20works%20with%20Linux%20and,with%20all%20Window%20Server%20platforms.>

upvoted 2 times

Buruguduystunstugudunstuy 1 year, 5 months ago

Selected Answer: B

A. Configure AWS Storage Gateway in volume gateway mode. Mount the volume to each Windows instance.

This option is incorrect because AWS Storage Gateway is not a file storage service. It is a hybrid storage service that allows you to store data in the cloud while maintaining low-latency access to frequently accessed data. It is designed to integrate with on-premises storage systems, not to provide file storage for Amazon EC2 instances.

B. Configure Amazon FSx for Windows File Server. Mount the Amazon FSx file system to each Windows instance.

This is the correct answer. Amazon FSx for Windows File Server is a fully managed file storage service that provides a native Windows file system that can be accessed over the SMB protocol. It is specifically designed for use with Windows-based applications, and it can be easily integrated with existing applications by mounting the file system to each EC2 instance.

upvoted 4 times

Buruguduystunstugudunstuy 1 year, 5 months ago

C. Configure a file system by using Amazon Elastic File System (Amazon EFS). Mount the EFS file system to each Windows instance.

This option is incorrect because Amazon EFS is a file storage service that is designed for use with Linux-based applications. It is not compatible with Windows-based applications, and it cannot be accessed over the SMB protocol.

D. Configure an Amazon Elastic Block Store (Amazon EBS) volume with the required size. Attach each EC2 instance to the volume. Mount the file system within the volume to each Windows instance.

This option is incorrect because Amazon EBS is a block storage service, not a file storage service. It is designed for storing raw block-level data that can be accessed by a single EC2 instance at a time. It is not designed for use as a shared file system that can be accessed by multiple instances.

upvoted 1 times

BENICE 1 year, 5 months ago

B - is correct

upvoted 1 times

career360guru 1 year, 5 months ago

Selected Answer: B

Option B

upvoted 1 times

Wpcorgan 1 year, 6 months ago

B is correct

upvoted 1 times

xua81376 1 year, 6 months ago

B FSx for windows

upvoted 2 times

BENICE 1 year, 6 months ago

B is correct option

upvoted 1 times

rjam 1 year, 6 months ago

Selected Answer: B

Amazon FSx for Windows File Server

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #187

Topic 1

A company is developing an ecommerce application that will consist of a load-balanced front end, a container-based application, and a relational database. A solutions architect needs to create a highly available solution that operates with as little manual intervention as possible.

Which solutions meet these requirements? (Choose two.)

- A. Create an Amazon RDS DB instance in Multi-AZ mode. **Most Voted**
- B. Create an Amazon RDS DB instance and one or more replicas in another Availability Zone.
- C. Create an Amazon EC2 instance-based Docker cluster to handle the dynamic application load.
- D. Create an Amazon Elastic Container Service (Amazon ECS) cluster with a Fargate launch type to handle the dynamic application load. **Most Voted**
- E. Create an Amazon Elastic Container Service (Amazon ECS) cluster with an Amazon EC2 launch type to handle the dynamic application load.

Correct Answer: AD

Community vote distribution

AD (100%)

Comments

techhb **Highly Voted** 1 year, 11 months ago

Selected Answer: AD

<https://containeronaws.com/introduction/ec2-or-aws-fargate/>

A.(O) multi-az <= 'little intervention'

B.(X) read replica <= Promoting a read replica to be a standalone DB instance

You can promote a read replica into a standalone DB instance. When you promote a read replica, the DB instance is rebooted before it becomes available.

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_ReadRepl.html

C.(X) use Amazon ECS instead of EC2-based docker for little human intervention

D.(O) Amazon ECS on AWS Fargate : AWS Fargate is a technology that you can use with Amazon ECS to run containers without having to manage servers or clusters of Amazon EC2 instances.

E.(X) EC2 launch type

The EC2 launch type can be used to run your containerized applications on Amazon EC2 instances that you register to your Amazon ECS cluster and manage yourself.

invited 14 times

upvoted 14 times

Danny23132412141_2312 Highly Voted 2 years ago

AD for sure

Link: <https://www.examttopics.com/discussions/amazon/view/43729-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 5 times

satyaamm Most Recent 4 months ago

Selected Answer: AD

RDS Multi AZ and AWS Fargate are the most suitable here.

upvoted 1 times

lostmagnet001 10 months, 2 weeks ago

Selected Answer: AD

Highly available application - Amazon RDS DB instance in Multi-AZ

little manual intervention - Fargate

upvoted 2 times

TariqKipkemei 1 year, 3 months ago

Selected Answer: AD

highly available application, little manual intervention = serverless = Amazon Elastic Container Service with Fargate and Amazon RDS DB instance in Multi-AZ mode

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: AD

The correct answers are A and D.

A) Creating an RDS DB instance in Multi-AZ mode provides automatic failover to a standby replica in another Availability Zone, providing high availability.

D) Using ECS Fargate removes the need to provision and manage EC2 instances, allowing the service to scale dynamically based on demand. ECS handles load balancing and availability out of the box.

upvoted 2 times

jkirancdev 1 year, 4 months ago

Selected Answer: AD

AD is the correct answer

upvoted 2 times

cookieMr 1 year, 5 months ago

Selected Answer: AD

A. Create an Amazon RDS DB instance in Multi-AZ mode. This ensures that the database is highly available with automatic failover to a standby replica in another Availability Zone.

D. Create an Amazon Elastic Container Service (Amazon ECS) cluster with a Fargate launch type to handle the dynamic application load. Fargate abstracts the underlying infrastructure, automatically scaling and managing the containers, making it a highly available and low-maintenance option.

Option B is not the best choice as it only creates replicas in another Availability Zone without the automatic failover capability provided by Multi-AZ mode.

Option C is not the best choice as managing a Docker cluster on EC2 instances requires more manual intervention compared to using the serverless capabilities of Fargate in option D.

Option E is not the best choice as it uses the EC2 launch type, which requires managing and scaling the EC2 instances manually. Fargate, as mentioned in option D, provides a more automated and scalable solution.

upvoted 4 times

studynoplay 1 year, 7 months ago

Selected Answer: AD

little manual intervention = Serverless

little manual intervention = Serverless

upvoted 2 times

career360guru 1 year, 12 months ago

Selected Answer: AD

Option A&D

upvoted 1 times

k1kavi1 1 year, 12 months ago

Selected Answer: AD

A and D

upvoted 2 times

Gabs90 2 years ago

Selected Answer: AD

A and D

upvoted 1 times

Wpcorgan 2 years ago

A and D

upvoted 1 times

BENICE 2 years ago

A and D are the options

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #188

Topic 1

A company uses Amazon S3 as its data lake. The company has a new partner that must use SFTP to upload data files. A solutions architect needs to implement a highly available SFTP solution that minimizes operational overhead.

Which solution will meet these requirements?

- A. Use AWS Transfer Family to configure an SFTP-enabled server with a publicly accessible endpoint. Choose the S3 data lake as the destination. **Most Voted**
- B. Use Amazon S3 File Gateway as an SFTP server. Expose the S3 File Gateway endpoint URL to the new partner. Share the S3 File Gateway endpoint with the new partner.
- C. Launch an Amazon EC2 instance in a private subnet in a VPC. Instruct the new partner to upload files to the EC2 instance by using a VPN. Run a cron job script, on the EC2 instance to upload files to the S3 data lake.
- D. Launch Amazon EC2 instances in a private subnet in a VPC. Place a Network Load Balancer (NLB) in front of the EC2 instances. Create an SFTP listener port for the NLB. Share the NLB hostname with the new partner. Run a cron job script on the EC2 instances to upload files to the S3 data lake.

Correct Answer: A

Community vote distribution

A (100%)

Comments

roxx529 **Highly Voted** 2 years ago

For Exam :
Whenever you see SFTP , FTP look for "Transfer" in options available
upvoted 65 times

LoXoL 1 year, 4 months ago

+ FTPS
upvoted 4 times

Chirantan **Highly Voted** 2 years, 5 months ago

Answer is A
AWS Transfer Family securely scales your recurring business-to-business file transfers to AWS Storage services using SFTP, FTPS

AWS Transfer Family securely scales your recurring business-to-business file transfers to AWS Storage services using SFTP, FTPS, FTP, and AS2 protocols.

<https://aws.amazon.com/aws-transfer-family/>

upvoted 17 times

oguzbeliren 1 year, 10 months ago

Answer A is not an answer because it requires more manual effort. While AWS Transfer Family simplifies the setup of an SFTP server, it still requires management and monitoring. This includes handling scaling, backups, patching, and other administrative tasks associated with managing an SFTP server.

upvoted 2 times

satyaamm Most Recent 4 months ago

Selected Answer: A

AWS Transfer Family is for FTP,SFTP,FTPS

upvoted 1 times

JohnZh 1 year, 2 months ago

Option A is the most suitable choice for implementing a highly available SFTP solution with minimal operational overhead in this scenario

upvoted 2 times

thewalker 1 year, 4 months ago

Selected Answer: A

The key advantages of AWS Transfer Family are:

It provides a fully managed file transfer service that eliminates the need to manage your own file transfer infrastructure. This reduces operational overhead.

It supports multiple protocols like SFTP, FTPS, FTP and AS2, allowing easy and secure exchange of data with business partners and customers.

File transfers happen directly into Amazon S3 buckets or Amazon EFS file systems, so the transferred data can be easily accessed by other AWS services for analytics, processing etc.

AWS Transfer Family maintains existing client-side configurations, so file transfer workflows remain unchanged for end users and partners.

It provides high availability and auto-scaling capabilities to handle varying transfer workloads.

upvoted 7 times

thewalker 1 year, 4 months ago

Storing transferred files in AWS allows using a broad range of services for compliance, archiving and deriving insights from the data.

AWS manages the file transfer infrastructure so you don't have to provision, operate and maintain file transfer servers.

For more details on AWS Transfer Family features, pricing and quotas, please refer to the documentation at

<https://aws.amazon.com/transfer-family>

upvoted 2 times

ansagr 1 year, 6 months ago

Selected Answer: A

Amazon S3 File Gateway, involves deploying an on-premises gateway that interfaces with S3. While it's a valid solution, it introduces a level of on-premises infrastructure that may require more operational management.

upvoted 2 times

JA2018 6 months, 3 weeks ago

just to add on, stem did not manage if the other party is accessing the owner's S3 from an on-prem site.

upvoted 1 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: A

AWS Transfer Family securely scales your recurring business-to-business file transfers to AWS Storage services using SFTP, FTPS, FTP, and AS2 protocols.

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

A is the correct answer.

AWS Transfer Family provides a fully managed SFTP service that can integrate directly with S3. It handles scaling, availability, and security automatically with minimal overhead.

upvoted 3 times

oguzbeliren 1 year, 10 months ago

AWS Transfer Family is a fully managed service that makes it easy to set up and manage secure file transfers. It provides a high-availability SFTP server that can be accessed from the public internet. However, this solution does not minimize operational overhead, as it requires the solutions architect to manage the SFTP server.

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: A

This solution provides a highly available SFTP solution without the need for manual management or operational overhead. AWS Transfer Family allows you to easily set up an SFTP server with authentication, authorization, and integration with S3 as the storage backend.

Option B is not the best choice as it suggests using Amazon S3 File Gateway, which is primarily used for file-based access to S3 storage over NFS or SMB protocols, not for SFTP access.

Option C is not the best choice as it requires manual management of an EC2 instance, VPN setup, and cron job script for uploading files, introducing operational overhead and potential complexity.

Option D is not the best choice as it also requires manual management of EC2 instances, Network Load Balancer, and cron job scripts for file uploads. It is more complex and involves additional components compared to the simpler and fully managed solution provided by AWS Transfer Family in option A.

upvoted 4 times

cookieMr 1 year, 11 months ago

This solution provides a highly available SFTP solution without the need for manual management or operational overhead. AWS Transfer Family allows you to easily set up an SFTP server with authentication, authorization, and integration with S3 as the storage backend.

Option B is not the best choice as it suggests using Amazon S3 File Gateway, which is primarily used for file-based access to S3 storage over NFS or SMB protocols, not for SFTP access.

Option C is not the best choice as it requires manual management of an EC2 instance, VPN setup, and cron job script for uploading files, introducing operational overhead and potential complexity.

Option D is not the best choice as it also requires manual management of EC2 instances, Network Load Balancer, and cron job scripts for file uploads. It is more complex and involves additional components compared to the simpler and fully managed solution provided by AWS Transfer Family in option A.

upvoted 4 times

cookieMr 1 year, 11 months ago

A is correct

upvoted 2 times

markw92 1 year, 12 months ago

I can't wrap my head around why the answer is D? this is so frustrating to see where i went wrong. I vote for A.

upvoted 3 times

studynoplay 2 years, 1 month ago

Selected Answer: A

minimizes operational overhead = Serverless

AWS Transfer Family is serverless

upvoted 2 times

Rahulbit34 2 years, 1 month ago

AWS Transfer Family is compatible for SFTP<FTPS<FTP. A is the answer

upvoted 2 times

kruasan 2 years, 1 month ago

Selected Answer: A

AWS Transfer Family is a fully managed AWS service that you can use to transfer files into and out of Amazon Simple Storage Service (Amazon S3) storage or Amazon Elastic File System (Amazon EFS) file systems over the following protocols:

Secure Shell (SSH) File Transfer Protocol (SFTP): version 3

File Transfer Protocol Secure (FTPS)

File Transfer Protocol (FTP)

Applicability Statement 2 (AS2)

upvoted 4 times

Oyz 2 years, 2 months ago

Selected Answer: A

A - is the correct answer.

upvoted 3 times

BENICE 2 years, 5 months ago

A -- is the option

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #189

Topic 1

A company needs to store contract documents. A contract lasts for 5 years. During the 5-year period, the company must ensure that the documents cannot be overwritten or deleted. The company needs to encrypt the documents at rest and rotate the encryption keys automatically every year.

Which combination of steps should a solutions architect take to meet these requirements with the LEAST operational overhead? (Choose two.)

A. Store the documents in Amazon S3. Use S3 Object Lock in governance mode.

B. Store the documents in Amazon S3. Use S3 Object Lock in compliance mode. **Most Voted**


C. Use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Configure key rotation.

D. Use server-side encryption with AWS Key Management Service (AWS KMS) customer managed keys. Configure key rotation. **Most Voted**

E. Use server-side encryption with AWS Key Management Service (AWS KMS) customer provided (imported) keys. Configure key rotation.

Correct Answer: BD

Community vote distribution



Comments

[Removed] **Highly Voted** 2 years, 6 months ago

Selected Answer: BD

Originally answered B and C due to least operational overhead. after research its bugging me that the s3 key rotation is determined based on AWS master Key rotation which cannot guarantee the key is rotated with in a 365 day period. stated as "varies" in the documentation. also its impossible to configure this in the console.

KMS-C is a tick box in the console to turn on annual key rotation but requires more operational overhead than SSE-S3.

C - will not guarantee the questions objectives but requires little overhead.

D - will guarantee the questions objective with more overhead.

upvoted 26 times

vadiminski_a 2 years, 5 months ago

I'd have to disagree on that. It states here that aws managed keys are rotated every year which is what the question asks:

to have to disagree on that states here that customer managed keys are rotated every year. Which is what the question asks.
<https://docs.aws.amazon.com/kms/latest/developerguide/rotate-keys.html> so C would be correct.
However, it also states that you cannot enable or disable rotation for aws managed keys which would again point towards D
upvoted 4 times

jdr75 2 years, 2 months ago

You can't use this link
<https://docs.aws.amazon.com/kms/latest/developerguide/rotate-keys.html>
to said that "sse-s3" rotates every year, cos' precisely that link refers to "KMS", that is covered with option D.
That the reason the solution is B+D.
upvoted 3 times

Dharmarajan 4 months, 1 week ago

The doc says:

"By default, when you enable automatic key rotation for a KMS key, AWS KMS generates new cryptographic material for the KMS key every year. You can also specify a custom rotation-period to define the number of days after you enable automatic key rotation that AWS KMS will rotate your key material, and the number of days between each automatic rotation thereafter. If you need to immediately initiate key material rotation, you can perform on-demand rotation, regardless of whether or not automatic key rotation is enabled. On-demand rotations do not change existing automatic rotation schedules."

Per this, the KMS does allow annual rotation and with an option to set frequency.
upvoted 1 times

pentium75 Highly Voted 1 year, 5 months ago

Selected Answer: BD

A - Governance mode allows exceptions
B - Yes
C - SSE-S3 rotates keys when AWS thinks is right, not when customer wants ("every year")
D - Yes
E - "customer provided (imported) keys" can obviously not be 'rotated automatically', the customer would have to provide/import new keys.
upvoted 9 times

celestial39 1 year, 4 months ago

KMS indeed rotates keys every year, but the reason why C is wrong is that the Amazon managed keys can't be configured to rotate or not.
REF: <https://docs.aws.amazon.com/kms/latest/developerguide/rotate-keys.html#rotate-keys-how-it-works>
upvoted 2 times

LoXoL 1 year, 4 months ago

Agree with pentium75
upvoted 3 times

Yak_Yeti Most Recent 2 months ago

Selected Answer: BC

SSE-S3 has the least operational overhead for encryption at rest with automatic key rotation because AWS fully manages the keys and their rotation (even if the schedule isn't user-configurable to be strictly annual).
However, if the strict annual rotation is a non-negotiable requirement, then you would have to accept the slightly higher (but still relatively low) operational overhead of managing a KMS customer managed key with automatic annual rotation enabled. Option D.
Given the wording of the original question prioritizing "least operational overhead," the intended answer likely leans towards the fully managed nature of SSE-S3's automatic rotation, even without a guaranteed annual cycle. If strict annual rotation was paramount, the question should have been phrased differently.
IMHO
upvoted 1 times

bignatov 9 months, 1 week ago

Selected Answer: BC

For me it is B and C. C because it does everything automatically and this is the requirement. (with LEAST operational overhead)
upvoted 1 times

tch 3 months ago

I think Amazon S3 managed keys (SSE-S3) are designed for and exclusively support encryption of objects within Amazon S3 storage only.....

upvoted 1 times

Americanman 9 months, 3 weeks ago

S3 Object Lock can help prevent Amazon S3 objects from being deleted or overwritten for a fixed amount of time or indefinitely.

In compliance mode, a protected object version can't be overwritten or deleted by any user, including the root user in your AWS account.

upvoted 2 times

ChymKuBoy 11 months, 3 weeks ago

Selected Answer: BD

BD for sure

upvoted 2 times

lofzee 1 year ago

Selected Answer: BD

basically what that pentium75 guy said - correct.

upvoted 2 times

sudohogan 1 year, 1 month ago

"Least operational overhead": C

upvoted 2 times

huangyou2003 1 year, 1 month ago

Selected Answer: BD

C- you don't have control over rotation schedule for SSE-S3

upvoted 2 times

Tralfalgarlaw 1 year, 1 month ago

Selected Answer: BD

B. Using S3 Object Lock in compliance mode ensures that the documents cannot be substituted or deleted during the specified retention period, which in this case is 5 years. This helps meet the requirement of ensuring the documents remain immutable for the duration of the contract.

D. Using server-side encryption with AWS Key Management Service (AWS KMS) customer managed keys allows for encryption of the documents at rest. Additionally, configuring key rotation for the customer managed keys ensures that the encryption keys are automatically rotated every year, meeting the requirement of rotating encryption keys automatically.

upvoted 3 times

MehulKapadia 1 year, 2 months ago

Selected Answer: BD

Answer: BD:

B: S3 Compliance Mode ensures no one can overwrite or delete the object.

D: Customer-managed KMS Key: (must be enabled) automatic every 1 year

Options not right:

A: Governance mode allows override and delete.

C: SSE-S3 customer do not have control on rotation of keys(Which is once a year in our requirement)

E: As per AWS Documentation, Customer Imported keys cannot be auto rotated.

upvoted 2 times

scar0909 1 year, 3 months ago

Selected Answer: BD

<https://docs.aws.amazon.com/kms/latest/developerguide/rotate-keys.html>

upvoted 1 times

thewalker 1 year, 4 months ago

mcmanus 1 year, 1 month ago

Selected Answer: BD

The best option to encrypt data at rest in Amazon S3 and rotate the keys every year is to use AWS KMS (Key Management Service).

With AWS KMS:

You can create a customer master key (CMK) and schedule automatic key rotation every year. This ensures the data is encrypted with a new key annually.

When storing objects in S3, you can choose server-side encryption with AWS KMS (SSE-KMS). This will encrypt the data with the CMK you created.

Even if the encrypted data is copied or transferred, it will remain encrypted since the keys are managed by KMS.

You have full control over the keys and can define IAM policies for key access.

AWS manages the encryption, key operations and auditing through integrated services like CloudTrail.

It provides an end-to-end encryption solution within AWS without needing to handle encryption/decryption yourself.

upvoted 1 times

omarshaban 1 year, 4 months ago

THIS WAS IN MY EXAM

upvoted 3 times

Mikado211 1 year, 5 months ago

File cannot be overwritten = s3 compliance mode

encryption AT REST = user-side encryption

upvoted 1 times

Mikado211 1 year, 5 months ago

so the correct answer is BD

upvoted 1 times

awsgeek75 1 year, 4 months ago

user side encryption?

upvoted 1 times

Mikado211 1 year, 5 months ago

Selected Answer: BD

File cannot be overwritten = compliance mode

Encryption AT REST = user-side encryption

upvoted 3 times

ale_brd_111 1 year, 6 months ago

Selected Answer: BD

Question might be outdated.

Amazon S3 now automatically applies server-side encryption with Amazon S3 managed keys (SSE-S3) as the default encryption for all buckets since January 5, 2023.

Additionally, it encrypts the key itself with another key that undergoes regular rotation, enhancing security.

Regarding key rotation, the document specifies that the key used to encrypt the S3 Encryption Key undergoes regular rotation.

However, it does not explicitly mention the rotation frequency or the ability to customize it.

Therefore, considering the requirement for key rotation and the lack of explicit details about rotation frequency, options B and D would be suitable choices.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #190

Topic 1


A company has a web application that is based on Java and PHP. The company plans to move the application from on premises to AWS. The company needs the ability to test new site features frequently. The company also needs a highly available and managed solution that requires minimum operational overhead.

Which solution will meet these requirements?

- A. Create an Amazon S3 bucket. Enable static web hosting on the S3 bucket. Upload the static content to the S3 bucket. Use AWS Lambda to process all dynamic content.
- B. Deploy the web application to an AWS Elastic Beanstalk environment. Use URL swapping to switch between multiple Elastic Beanstalk environments for feature testing. **Most Voted**
- C. Deploy the web application to Amazon EC2 instances that are configured with Java and PHP. Use Auto Scaling groups and an Application Load Balancer to manage the website's availability.
- D. Containerize the web application. Deploy the web application to Amazon EC2 instances. Use the AWS Load Balancer Controller to dynamically route traffic between containers that contain the new site features for testing.

Correct Answer: B

Community vote distribution



Comments

cookieMr **Highly Voted** 1 year, 5 months ago

Selected Answer: B

B. Provides a highly available and managed solution with minimum operational overhead. By deploying the web application to EBS, the infrastructure and platform management are abstracted, allowing easy deployment and scalability. With URL swapping, different environments can be created for testing new site features, and traffic can be routed between these environments without any downtime.

A. Suggests using S3 for static content hosting and Lambda for dynamic content. While it offers simplicity for static content, it does not provide the necessary flexibility and dynamic functionality required by a Java and PHP-based web application.

C. Involves manual management of EC2, ASG, and ELB, which requires more operational overhead and may not provide the desired level of availability and ease of testing.

D. Introduces containerization, which adds complexity and operational overhead for managing containers and infrastructure.

increases complexity, which adds complexity and operational overhead for managing containers and infrastructure, making it less suitable for a requirement of minimum operational overhead.

upvoted 16 times

Shasha1 Highly Voted 2 years ago

B
Elastic Beanstalk is a fully managed service that makes it easy to deploy and run applications in the AWS; To enable frequent testing of new site features, you can use URL swapping to switch between multiple Elastic Beanstalk environments.

upvoted 14 times

oguzbeliren 1 year, 4 months ago

The correct answer is D.

AWS Elastic Beanstalk is a service that makes it easy to deploy and manage web applications in the AWS cloud. However, it is not a good solution for testing new site features frequently, as it can be difficult to switch between multiple Elastic Beanstalk environments.

upvoted 3 times

satyaamm Most Recent 4 months ago

Selected Answer: B

AWS Elastic Beanstalk is most suitable for deploying and managing web applications.

upvoted 1 times

sudohogan 7 months, 1 week ago

Containers allow you test your app in isolated environments, therefore D is the correct option.

upvoted 1 times

hro 8 months, 2 weeks ago

B - Because AWS Elastic Beanstalk performs an in-place update when you update your application versions, your application might become unavailable to users for a short period of time. To avoid this, perform a blue/green deployment. To do this, deploy the new version to a separate environment, and then swap the CNAMEs of the two environments to redirect traffic to the new version instantly.

upvoted 3 times

reviewmine 9 months, 4 weeks ago

Selected Answer: B

Elastic Beanstalk can test Blue/Green deployment. Switching Dev to prod/ prod to dev easily.

upvoted 3 times

awsgeek75 11 months, 2 weeks ago

Selected Answer: B

A and C are not allowing for feature testing.

B and D allow feature testing. D requires overhead of containerisation as well as the LB controller to selectively chose containers for features (assuming on how this might be implemented). EBS allows switching between environment like A/B testing but on whole site. Expensive but cost is not a concern for this question.

upvoted 3 times

master9 11 months, 3 weeks ago

Selected Answer: D

WS Elastic Beanstalk supports multiple environments, but each environment can only run one platform at a time. A platform is a combination of an operating system, runtime, and web server, and in this case, Java and PHP would be considered different platforms.

So, if you want to use both Java and PHP, you would need to create two separate environments, one for each. You can then link these environments together using AWS services like Route 53 for routing traffic, or use an Application Load Balancer to distribute incoming traffic between the two environments.

upvoted 1 times

ale_brd_111 1 year ago

Selected Answer: B

Option B (AWS Elastic Beanstalk): Elastic Beanstalk is a fully managed service that makes it easy to deploy and run applications

Open to the Elastic Beanstalk, Elastic Beanstalk is a fully managed service that makes it easy to deploy and run applications in multiple languages (including Java and PHP) with MINIMAL OPERATION OVERHEAD. It abstracts the infrastructure management, allowing you to focus on your application. URL swapping in Elastic Beanstalk allows you to easily switch between different environments, making it convenient for testing new features.

upvoted 2 times

Po_chih 1 year, 2 months ago

Selected Answer: B

B
Elastic Beanstalk is a fully managed service that makes it easy to deploy and run applications in the AWS; To enable frequent testing of new site features, you can use URL swapping to switch between multiple Elastic Beanstalk environments.
https://docs.aws.amazon.com/zh_tw/whitepapers/latest/blue-green-deployments/swap-the-environment-of-an-elastic-beanstalk-application.html

upvoted 4 times

Po_chih 1 year, 2 months ago

Selected Answer: B

B
Elastic Beanstalk is a fully managed service that makes it easy to deploy and run applications in the AWS; To enable frequent testing of new site features, you can use URL swapping to switch between multiple Elastic Beanstalk environments.
https://docs.aws.amazon.com/zh_tw/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html

upvoted 2 times

TariqKipkemei 1 year, 3 months ago

Selected Answer: B

AWS Elastic Beanstalk URL swapping is the main ask of this question.

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

B is the correct answer.

Using AWS Elastic Beanstalk provides a fully managed platform to deploy the web application. Elastic Beanstalk will handle provisioning EC2 instances, load balancing, auto scaling, and application health monitoring.

Elastic Beanstalk's ability to support multiple environments and swap URLs allows easy testing of new features before swapping into production. This requires minimal overhead compared to managing infrastructure directly.

upvoted 4 times

oguzbeliren 1 year, 4 months ago

The correct answer is D.

AWS Elastic Beanstalk is a service that makes it easy to deploy and manage web applications in the AWS cloud. However, it is not a good solution for testing new site features frequently, as it can be difficult to switch between multiple Elastic Beanstalk environments.

upvoted 1 times

Abrar2022 1 year, 6 months ago

S3 is for hosting static websites not dynamic websites or applications
Beanstalk will take care of this.

upvoted 1 times

kruasan 1 year, 7 months ago

Selected Answer: B

Frequent feature testing -
- Multiple Elastic Beanstalk environments can be created easily for development, testing and production use cases.
- Traffic can be routed between environments for A/B testing and feature iteration using simple URL swapping techniques. No complex routing rules or infrastructure changes required.

upvoted 1 times

ashu089 1 year, 8 months ago

who needs discussion in the era the of chatGPT

upvoted 3 times

aadityaravi8 1 year, 5 months ago

chatGPT always change its answer. just say wrong answer, he will come up with new answer each time with justification. chatGPT is not trusted at all.

upvoted 4 times

baku98 1 year ago

In the era of ChatGPT, individuals across education, business, content creation, healthcare, programming, language learning, innovation, and mental health benefit from discussions. Students, professionals, writers, developers, language learners, innovators, and those seeking support find ChatGPT valuable for learning, problem-solving, creative endeavors, and companionship. It serves as a versatile tool for information, collaboration, and engagement across diverse domains, enhancing communication in an accessible and interactive manner.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #191

Topic 1

A company has an ordering application that stores customer information in Amazon RDS for MySQL. During regular business hours, employees run one-time queries for reporting purposes. Timeouts are occurring during order processing because the reporting queries are taking a long time to run. The company needs to eliminate the timeouts without preventing employees from performing queries.

What should a solutions architect do to meet these requirements?

- A. Create a read replica. Move reporting queries to the read replica. **Most Voted**
- B. Create a read replica. Distribute the ordering application to the primary DB instance and the read replica.
- C. Migrate the ordering application to Amazon DynamoDB with on-demand capacity.
- D. Schedule the reporting queries for non-peak hours.

Correct Answer: A

Community vote distribution

A (100%)

Comments

BENICE **Highly Voted** 2 years, 5 months ago

A is correct answer. This was in my exam
upvoted 28 times

Grace83 2 years, 2 months ago

Did these questions help with your exam?
upvoted 5 times

zied007 **Most Recent** 8 months, 3 weeks ago

Selected Answer: A

The best solution to address the timeouts and eliminate the impact of long-running reporting queries without disrupting order processing is: A
upvoted 2 times

PR5577 10 months, 3 weeks ago

Selected Answer: A

Reporting queries will point to read replica. Application will still point to primary db for write / read operations.
upvoted 2 times

bakstorage00001 1 year, 2 months ago

Selected Answer: A

The correct answer to this question is A: Create a read replica. Move reporting queries to the read replica. This solution is designed to alleviate the load on the primary database used by the ordering application. By offloading the reporting queries to a read replica, the primary instance is freed up to handle operational transactions like order processing without contention from the resource-intensive reporting queries. This should effectively reduce or eliminate the timeouts currently experienced during order processing.

upvoted 3 times

lostmagnet001 1 year, 4 months ago

Selected Answer: A

create the replica and all the report queries get data from that read replica.

upvoted 2 times

truongtx8 1 year, 4 months ago

Selected Answer: A

B incorrect because ordering application needs to write data to the DB.

upvoted 3 times

Ruffyt 1 year, 7 months ago

A. By moving the reporting queries to the read replica, the primary DB instance used for order processing is not affected by the long-running reporting queries. This helps eliminate timeouts during order processing while allowing employees to perform their queries without impacting the application's performance.

B. While this can provide some level of load distribution, it does not specifically address the issue of timeouts caused by reporting queries during order processing.

C. While DynamoDB offers scalability and performance benefits, it may require significant changes to the application's data model and querying approach.

D. While this approach can help alleviate the impact on order processing, it does not address the requirement of eliminating timeouts without preventing employees from performing queries.

upvoted 3 times

David_Ang 1 year, 7 months ago

Selected Answer: A

"A" is correct because it does not cause problems in the primary DB

upvoted 2 times

TariqKipkemei 1 year, 9 months ago

Selected Answer: A

reports = read replica

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

A is the correct answer.

Creating an RDS MySQL read replica will allow the reporting queries to be isolated and run without affecting performance of the primary ordering application.

Read replicas allow read-only workloads to be scaled out while eliminating contention with the primary write workload.

upvoted 3 times

james2033 1 year, 10 months ago

Selected Answer: A

Question keyword "regular business hours" made D is incorrect.

C migrate to Amazon DynamoDB (No-SQL) is meaningless, remove C.

Answer B, create a "read replica", it is ok, but "ordering application pointed to read replica" is incorrect.

A is correct answer. Easy question.

upvoted 4 times

sickcow 1 year, 11 months ago

Selected Answer: A

A sounds right

upvoted 2 times

rauldevilla 1 year, 11 months ago

Selected Answer: A

Using the primary instance continues with the problem

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: A

A. By moving the reporting queries to the read replica, the primary DB instance used for order processing is not affected by the long-running reporting queries. This helps eliminate timeouts during order processing while allowing employees to perform their queries without impacting the application's performance.

B. While this can provide some level of load distribution, it does not specifically address the issue of timeouts caused by reporting queries during order processing.

C. While DynamoDB offers scalability and performance benefits, it may require significant changes to the application's data model and querying approach.

D. While this approach can help alleviate the impact on order processing, it does not address the requirement of eliminating timeouts without preventing employees from performing queries.

upvoted 4 times

steev 2 years ago

Selected Answer: A

correct

upvoted 2 times

cheese929 2 years, 1 month ago

Selected Answer: A

A is correct.

upvoted 2 times

kruasan 2 years, 1 month ago

Selected Answer: A

Creating a read replica allows the company to offload the reporting queries to a separate database instance, reducing the load on the primary database used for order processing. By moving the reporting queries to the read replica, the ordering application running on the primary DB instance can continue to process orders without timeouts due to the long-running reporting queries.

Option B is not a good solution because distributing the ordering application to the primary DB instance and the read replica does not address the issue of long-running reporting queries causing timeouts during order processing.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #192

Topic 1

A hospital wants to create digital copies for its large collection of historical written records. The hospital will continue to add hundreds of new documents each day. The hospital's data team will scan the documents and will upload the documents to the AWS Cloud.

A solutions architect must implement a solution to analyze the documents, extract the medical information, and store the documents so that an application can run SQL queries on the data. The solution must maximize scalability and operational efficiency.

Which combination of steps should the solutions architect take to meet these requirements? (Choose two.)

A. Write the document information to an Amazon EC2 instance that runs a MySQL database.

B. Write the document information to an Amazon S3 bucket. Use Amazon Athena to query the data. **Most Voted**

C. Create an Auto Scaling group of Amazon EC2 instances to run a custom application that processes the scanned files and extracts the medical information.

D. Create an AWS Lambda function that runs when new documents are uploaded. Use Amazon Rekognition to convert the documents to raw text. Use Amazon Transcribe Medical to detect and extract relevant medical information from the text.

E. Create an AWS Lambda function that runs when new documents are uploaded. Use Amazon Textract to convert the documents to raw text. Use Amazon Comprehend Medical to detect and extract relevant medical information from the text.

Most Voted**Correct Answer: BE***Community vote distribution*

BE (100%)

Comments**KADSM** **Highly Voted** 2 years ago

B and E are correct. Textract to extract text from files. Rekognition can also be used for text detection but after Rekognition - it's mentioned that Transcribe is used. Transcribe is used for Speech to Text. So that option D may not be valid.
upvoted 14 times

Kapello10 **Highly Voted** 2 years ago

B and E is the correct ans

B > Store documents on S3 and use Athena to query >

E > Use Textract to extract text from files and not Rekognition. N.B Rekognition is for image identification

upvoted 5 times

satyaamm Most Recent 3 months, 4 weeks ago

Selected Answer: BE

Athena for SQL queries and Lambda for scalability.

upvoted 1 times

LoXoL 11 months ago

Selected Answer: BE

no brainer: B,E

upvoted 4 times

awsgeek75 11 months, 2 weeks ago

Selected Answer: BE

E: Amazon Textract & Amazon Comprehend Medical obviously do the job with least operational overhead. D can do this but it will be extra work and overhead.

B for running SQL queries on S3 bucket directly without extra overhead.

upvoted 4 times

Ruffyt 1 year, 1 month ago

Write the document information to an Amazon S3 bucket. Use Amazon Athena to query the data.

Create an AWS Lambda function that runs when new documents are uploaded. Use Amazon Textract to convert the documents to raw text. Use Amazon Comprehend Medical to detect and extract relevant medical information from the text.

upvoted 2 times

David_Ang 1 year, 1 month ago

Selected Answer: BE

another mistake from the admin, should correct this one, because we all agree

upvoted 4 times

vijaykamal 1 year, 2 months ago

Answer - BE

Option D mentions using Amazon Rekognition and Amazon Transcribe Medical, which are primarily designed for image and audio analysis, respectively. While they can be part of a document processing pipeline, Amazon Textract and Amazon Comprehend Medical are more suitable for extracting structured information from documents, making option E a better choice.

upvoted 3 times

TariqKipkemei 1 year, 3 months ago

Selected Answer: BE

Write the document information to an Amazon S3 bucket. Use Amazon Athena to query the data.

Create an AWS Lambda function that runs when new documents are uploaded. Use Amazon Textract to convert the documents to raw text. Use Amazon Comprehend Medical to detect and extract relevant medical information from the text.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: BE

B and E are the correct answers.

B is correct because storing the scanned documents in Amazon S3 provides highly scalable and durable storage. Amazon Athena allows running SQL queries directly against the data in S3 without needing to load the data into a database.

E is correct because using Lambda functions triggered by uploads provides a serverless approach to automatically process each document. Amazon Textract and Comprehend Medical can extract text and medical information without needing to manage server

upvoted 5 times

james2033 1 year, 4 months ago

Selected Answer: BE

Amazon Comprehend Medical for image reading
<https://aws.amazon.com/comprehend/medical/> .

Amazon Transcribe Medical for speech audio. Remove D. Keep E.

A is meaningless, remove A (EC2).

B use Amazon S3, Athena for querying, keep B.

Conclusion combination B and E are correct answers.

upvoted 3 times

MNotABot 1 year, 5 months ago

AC wrong as involve EC2. Either one of DE are correct so that makes B correct. Now E is obvious answer if we have read AWS FAQs

upvoted 2 times

animefan1 1 year, 5 months ago

Selected Answer: BE

Textract to extract the content and Athena to run sql queries on S3 data

upvoted 2 times

sickcow 1 year, 5 months ago

Selected Answer: BE

From a DE/ML perspective Lambda + Textract + S3 + Athena is the best way to go

upvoted 2 times

rauldevilla 1 year, 5 months ago

Selected Answer: BE

Transcribe is used. Transcribe is used for Speech to Text

upvoted 2 times

cookieMr 1 year, 5 months ago

Selected Answer: BE

B is correct because it suggests writing the document information to an Amazon S3 bucket, which provides scalable and durable object storage. Using Amazon Athena, the data can be queried using SQL, enabling efficient analysis.

E is correct because it involves creating an AWS Lambda function triggered by new document uploads. Amazon Textract is used to convert the documents to raw text, and Amazon Comprehend Medical extracts relevant medical information from the text.

A is incorrect because writing the document information to an Amazon EC2 instance with a MySQL database is not a scalable or efficient solution for analysis.

C is incorrect because creating an Auto Scaling group of Amazon EC2 instances for processing scanned files and extracting information would introduce unnecessary complexity and management overhead.

D is incorrect because using an EC2 instance with a MySQL database for storing document information is not the optimal solution for scalability and efficient analysis.

upvoted 3 times

AlankarJ 1 year, 6 months ago

It states in the question that the written documents are scanned. They are converted into images after being scanned. Rekognition would be best to analyse images.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #193

Topic 1

A company is running a batch application on Amazon EC2 instances. The application consists of a backend with multiple Amazon RDS databases. The application is causing a high number of reads on the databases. A solutions architect must reduce the number of database reads while ensuring high availability.

What should the solutions architect do to meet this requirement?

A. Add Amazon RDS read replicas. **Most Voted**


B. Use Amazon ElastiCache for Redis.

C. Use Amazon Route 53 DNS caching

D. Use Amazon ElastiCache for Memcached.

Correct Answer: A

Community vote distribution

**Comments**

leonn **Highly Voted** 2 years, 6 months ago

Selected Answer: B

Use ElastiCache to reduce reading and choose redis to ensure high availability.
upvoted 42 times

Lalo 2 years, 3 months ago

Where is the high availability when the database fails and the cache time runs out?
The answer is a.
upvoted 26 times

mandragon 2 years ago

Elasticache for Redis ensures high availability by using read replicas and Multi AZ with failover. It is also faster since it uses cache.

<https://docs.aws.amazon.com/AmazonElastiCache/latest/red-ug/AutoFailover.html>
upvoted 5 times

LoXoL 1 year, 4 months ago

Right here: Redis has Multi AZ with Auto-Failover (in addition it got Read Replicas for HA)
upvoted 1 times

babayomi 8 months, 2 weeks ago

THE FOCUS HERE IS TO REDUCE READ OPERATION FROM THE READ REPLICAS. CATCHING IS THE SOLUTION AND REDIS IS MOST SUITABLE
upvoted 2 times

Mia2009687 1 year, 11 months ago

They run multiple databases
upvoted 1 times

JoeGuan 1 year, 9 months ago

Caching Frequently Accessed Data: ElastiCache allows you to store frequently accessed or computationally expensive data in-memory within the cache nodes. This means that when an application requests data, ElastiCache can provide the data directly from the cache without having to query the RDS database. This reduces the number of reads on the RDS database because the data is retrieved from the faster in-memory cache.
upvoted 6 times

channn **Highly Voted** 2 years, 2 months ago

Selected Answer: A

A vs B:

A: reduce the number of database reads on main + high availability provide

B: only reduce the number of DB reads

so A wins

upvoted 32 times

pentium75 1 year, 5 months ago

Why "only reduce the number of DB reads"? This is exactly what is asked.

Elasticache for Redis can be HA (contrary to Elasticache for Memcached).

upvoted 6 times

LoXoL 1 year, 4 months ago

pentium75 is right.

upvoted 2 times

network_enthusiast **Most Recent** 1 month, 4 weeks ago

Selected Answer: B

B is the correct answer

upvoted 1 times

kg508 3 months, 1 week ago

Selected Answer: B

I voted B because it was mentioned that the number of database reads should be reduced

If I had to reduce the database read load, I would have done it to A

upvoted 1 times

satyaammm 3 months, 4 weeks ago

Selected Answer: A

RDS read replicas provide HA unlike Redis

upvoted 1 times

FlyingHawk 4 months, 1 week ago

Selected Answer: B

Read replicas distribute read traffic but do not reduce the total number of reads. They scale read capacity but still require every read to hit a database instance.

upvoted 1 times

Qu_B 8 months, 1 week ago

This question is stupid! Both answers can be correct!

upvoted 4 times

Omariox 8 months, 1 week ago

Selected Answer: A

for high availability, the answer should be A

upvoted 4 times

jaradat02 10 months, 2 weeks ago

Selected Answer: B

RDS read replicas are mainly used for higher performance, and not for higher availability, that's why I chose ElastiCache using Redis.

upvoted 2 times

Aws_aspr 10 months, 3 weeks ago

Selected Answer: B

I Vote B since question is asking to reduce the database reads. Using Read replicas offloads the reading operation from main db but doesn't reduce it.

upvoted 2 times

freedafeng 11 months ago

B is not correct. the effectiveness of caching depends on the application. it's not always a good solution.

upvoted 1 times

jatric 11 months, 1 week ago

Selected Answer: A

why not read replica when you get read performant with HA?

upvoted 4 times

shil_31 1 year ago

Selected Answer: A

Options B, C, and D are not directly related to ensuring high availability

upvoted 4 times

lofzee 1 year ago

Selected Answer: B

I think it is B:

<https://aws.amazon.com/getting-started/hands-on/boosting-mysql-database-performance-with-amazon-elasticache-for-redis/>

Although the question just says "RDS Database", it doesn't specify what type of DB.

upvoted 1 times

sudohogan 1 year, 1 month ago

Elasticahe is NOSQL, A is the Answer!

upvoted 4 times

ManikRoy 1 year, 1 month ago

Selected Answer: B

Pretty sure answer is option B.

You have to use a caching to 'reduce' database reads so read replica is out of option.

The question mentions High availability so Redis is preferable than MemCached.

upvoted 3 times

soufiyane 1 year, 2 months ago

Selected Answer: B

the question is clearly says reduce reads request, so b is the naswer
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #194

Topic 1

A company needs to run a critical application on AWS. The company needs to use Amazon EC2 for the application's database. The database must be highly available and must fail over automatically if a disruptive event occurs.

Which solution will meet these requirements?

- A. Launch two EC2 instances, each in a different Availability Zone in the same AWS Region. Install the database on both EC2 instances. Configure the EC2 instances as a cluster. Set up database replication. **Most Voted**
- B. Launch an EC2 instance in an Availability Zone. Install the database on the EC2 instance. Use an Amazon Machine Image (AMI) to back up the data. Use AWS CloudFormation to automate provisioning of the EC2 instance if a disruptive event occurs.
- C. Launch two EC2 instances, each in a different AWS Region. Install the database on both EC2 instances. Set up database replication. Fail over the database to a second Region.
- D. Launch an EC2 instance in an Availability Zone. Install the database on the EC2 instance. Use an Amazon Machine Image (AMI) to back up the data. Use EC2 automatic recovery to recover the instance if a disruptive event occurs.

Correct Answer: A*Community vote distribution***Comments****Gil80** **Highly Voted** 2 years, 6 months ago**Selected Answer: A**

Changing my vote to A. After reviewing a Udemy course of SAA-C03, it seems that A (multi-AZ and Clusters) is sufficient for HA.
upvoted 39 times

berks 2 years, 5 months ago

what number of class ?
upvoted 5 times

AAAWrekng 1 year, 7 months ago

Here AWS defines HA, and uses the word cluster - AWS has several methods for achieving HA through both approaches, such

as through a scalable, load balanced cluster or assuming an active-standby pair. -

<https://docs.aws.amazon.com/whitepapers/latest/real-time-communication-on-aws/high-availability-and-scalability-on-aws.html>

upvoted 2 times

Gil80 Highly Voted 2 years, 6 months ago

Selected Answer: C

The question states that it is a critical app and it has to be HA. A could be the answer, but it's in the same AZ, so if the entire region fails, it doesn't cater for the HA requirement.

However, the likelihood of a failure in two different regions at the same time is 0. Therefore, to me it seems that C is the better option to cater for HA requirement.

In addition, C does state like A that the DB app is installed on an EC2 instance.

upvoted 29 times

Burrito69 1 year, 2 months ago

Option C proposes launching two EC2 instances in different AWS Regions and setting up database replication, with failover to a second Region. While this solution does provide geographic redundancy, it may introduce higher latency due to cross-region communication and data replication. Additionally, failover to a different Region typically involves more complex configurations and longer recovery times compared to failover within the same Region.

While Option C may offer a level of redundancy, it might not provide the same level of responsiveness and automatic failover capabilities as Option A, which leverages Availability Zones within the same Region. In scenarios where low latency and rapid failover are critical, Option A is often preferred. However, if geographic redundancy is a top priority and the potential trade-offs in latency and failover time are acceptable, Option C could still be a viable solution.

upvoted 7 times

javitech83 2 years, 6 months ago

but for C you need communication between the two VPC, which increase the complexity. With a should be enough for HA

upvoted 8 times

Steve_4542636 2 years, 3 months ago

The question doesn't ask which option is the most HA. It asks what meets the requirements.

upvoted 7 times

aussiehoa 2 years, 1 month ago

Design for region failure? may as well design for AWS failure and put replica in GCP and Azure :v

upvoted 12 times

jjcode 1 year, 6 months ago

this is reading comprehension exam not a practical exam.

upvoted 3 times

Kp88 1 year, 10 months ago

And on-prem in multiple DCs and one in mars too :D

upvoted 10 times

cyber_bedouin 1 year, 8 months ago

yep lol, even in the other questions, for HA you can use Multi-AZ

upvoted 3 times

bora4motion Most Recent 1 month, 2 weeks ago

Selected Answer: A

I don't understand people who voted for C really. A!

upvoted 1 times

SirDNS 2 months, 3 weeks ago

Selected Answer: A

Failover between Regions is typically slower than within a Region.

upvoted 1 times

Anastesas 3 months, 2 weeks ago

Selected Answer: A

Cross-region setups are primarily used for disaster recovery rather than high availability due to higher latency and complexity. T

upvoted 1 times

bignatov 8 months, 4 weeks ago

Selected Answer: A

I am voting for A.

upvoted 2 times

1e22522 10 months, 2 weeks ago

Selected Answer: A

It's a due to being in the same region different AZ for latency purposes.

upvoted 2 times

jatric 11 months, 1 week ago

Selected Answer: A

A provides HA with 2 EC2 in two AZ with database replication

upvoted 3 times

ChymKuBoy 11 months, 3 weeks ago

C for sure

upvoted 1 times

Duckydoo 11 months, 3 weeks ago

What does "Configure the EC2 instances as a cluster" mean? The only "EC2 cluster" that I am aware of is a "cluster placement group". If that's the case, then all EC2 instances in that cluster must be in the same AZ. So option A would be invalid then.

upvoted 2 times

shil_31 1 year ago

Selected Answer: A

Option B uses an AMI for backup and CloudFormation for automation, but doesn't provide high availability or automatic failover.

Option C launches instances in different regions, which may not be necessary and may increase costs.

Option D uses EC2 automatic recovery, which can recover an instance, but doesn't provide high availability or automatic failover.

upvoted 2 times

k_ 1 year ago

Selected Answer: C

The term "disruptive event" implies it requires DR, HA is not sufficient.

upvoted 3 times

Solomon2001 1 year, 1 month ago

Selected Answer: A

Explanation:

Option C provides a solution that ensures high availability by deploying EC2 instances in different AWS Regions. By setting up database replication and failing over the database to a second Region, you ensure automatic failover if a disruptive event occurs in one Region.

Options A and B focus on high availability within a single AWS Region but don't address automatic failover to a different Region in case of a disruptive event.

Option D uses EC2 automatic recovery, but it doesn't provide a solution for automatic failover to a different Region, which is necessary for ensuring high availability in case of a Region-wide failure.

upvoted 2 times

hardy1234567 1 year, 1 month ago

C - I am a little be wonder reading same explanation, because exist a vary big difference beetwen instance cluster and dababase cluster.

upvoted 1 times

Uzbekistan 1 year, 2 months ago

Selected Answer: A

Option A suggests deploying two EC2 instances, each in a different Availability Zone within the same AWS Region. This ensures high availability by distributing the instances across multiple physically isolated locations. By installing the database on both EC2 instances and configuring them as a cluster, you create a highly available database setup where one instance can seamlessly take over if the other instance fails.

Additionally, setting up database replication between the instances ensures data consistency and redundancy. If one instance fails, the other instance can continue serving requests without interruption.

upvoted 4 times

derekz 1 year, 3 months ago

Selected Answer: A

A is for HA. D is for DR

upvoted 2 times

MrPCarrot 1 year, 4 months ago

Perfect Answer is A

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #195

Topic 1


A company's order system sends requests from clients to Amazon EC2 instances. The EC2 instances process the orders and then store the orders in a database on Amazon RDS. Users report that they must reprocess orders when the system fails. The company wants a resilient solution that can process orders automatically if a system outage occurs.

What should a solutions architect do to meet these requirements?

- A. Move the EC2 instances into an Auto Scaling group. Create an Amazon EventBridge (Amazon CloudWatch Events) rule to target an Amazon Elastic Container Service (Amazon ECS) task.
- B. Move the EC2 instances into an Auto Scaling group behind an Application Load Balancer (ALB). Update the order system to send messages to the ALB endpoint.
- C. Move the EC2 instances into an Auto Scaling group. Configure the order system to send messages to an Amazon Simple Queue Service (Amazon SQS) queue. Configure the EC2 instances to consume messages from the queue. **Most Voted**
- D. Create an Amazon Simple Notification Service (Amazon SNS) topic. Create an AWS Lambda function, and subscribe the function to the SNS topic. Configure the order system to send messages to the SNS topic. Send a command to the EC2 instances to process the messages by using AWS Systems Manager Run Command.

Correct Answer: C

Community vote distribution



Comments

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: C

The key reasons are:

Using an Auto Scaling group ensures the EC2 instances that process orders are highly available and scalable. With SQS, the orders are decoupled from the instances that process them via asynchronous queuing. If instances fail or go down, the orders remain in the queue until new instances can pick them up. This provides automated resilience.

Any failed processing can retry by resending messages back to the queue

upvoted 16 times

awsgeek75 **Highly Voted** 11 months, 2 weeks ago

Selected Answer: C

A uses ECS tasks for something which makes no sense.
B does not solve the reliable processing of orders
C SQS for sending a message and processing it reliable
D is like reinventing SQS with SNS and Lambda mumbo jumbo
upvoted 5 times

satyaamm **Most Recent** 3 months, 4 weeks ago

Selected Answer: C

SQS queues are the most suitable here.
upvoted 1 times

jjcode 1 year ago

How does SNS capture the requests after the application fails? Those messages are ephemeral by nature and will not hold the data like SQS would. In theory one could create a subscription based service using SNS to stream the data to a service that could store the request, but why...
upvoted 2 times

pentium75 11 months, 3 weeks ago

That's one of the reasons why D is wrong (not to mention the "Systems Manager Run Command" nonsense).
upvoted 3 times

awsgeek75 11 months, 2 weeks ago

I stopped reading option D after SNS and Lambda.... it was sounding nonsense. SQS is default reliability delivery system for me.
upvoted 2 times

pavospam 1 year ago

Selected Answer: C

it's C... 4 answers wrong I have found
upvoted 1 times

Ruffyit 1 year, 1 month ago

C.
Option D suggests using Amazon SNS and AWS Lambda, which can be part of an event-driven architecture but may not be the best fit for ensuring the automatic processing of orders during system outages. It relies on an additional AWS Systems Manager Run Command step, which adds complexity and may not be as reliable as using SQS for queuing messages.
upvoted 2 times

David_Ang 1 year, 1 month ago

Selected Answer: C

"C" because they need to store the request and then be process by the system if it fails, SNS does not have that capacity.
another mistake from the admin
upvoted 2 times

vijaykamal 1 year, 2 months ago

Selected Answer: C

Option D suggests using Amazon SNS and AWS Lambda, which can be part of an event-driven architecture but may not be the best fit for ensuring the automatic processing of orders during system outages. It relies on an additional AWS Systems Manager Run Command step, which adds complexity and may not be as reliable as using SQS for queuing messages.
upvoted 2 times

TariqKipkemei 1 year, 3 months ago

Selected Answer: C

Move the EC2 instances into an Auto Scaling group. Configure the order system to send messages to an Amazon Simple Queue Service (Amazon SQS) queue. Configure the EC2 instances to consume messages from the queue.
upvoted 2 times

SumCloud 1 year, 3 months ago

Guru4Cloud 1 year, 5 months ago

Selected Answer: C

C is the correct answer.

Using an Auto Scaling group with EC2 instances behind a load balancer provides high availability and scalability.

Sending the orders to an SQS queue decouples the ordering system from the processing system. The EC2 instances can poll the queue for new orders and process them even during an outage. Any failed orders will go back to the queue for reprocessing.

upvoted 2 times

cookieMr 1 year, 5 months ago

Selected Answer: C

By moving the EC2 into an ASG and configuring them to consume messages from an SQS, the system can decouple the order processing from the order system itself. This allows the system to handle failures and automatically process orders even if the order system or EC2 experience outages.

A. Using an ASG with an EventBridge rule targeting an ECS task does not provide the necessary decoupling and message queueing for automatic order processing during outages.

B. Moving the EC2 instances into an ASG behind an ALB does not address the need for message queueing and automatic processing during outages.

D. Using SNS and Lambda can provide notifications and orchestration capabilities, but it does not provide the necessary message queueing and consumption for automatic order processing during outages. Additionally, using Systems Manager Run Command to send commands for order processing adds complexity and does not provide the desired level of automation.

upvoted 4 times

pisica134 1 year, 5 months ago

D is so unnecessary this confuses people

upvoted 2 times

cookieMr 1 year, 5 months ago

Thx Allmightly for voting system! Answers provided by the site (and not by community) are 20% wrong.

upvoted 4 times

markw92 1 year, 5 months ago

The answer D is so complex and unnecessary. Why moderator is not providing an explanation of answers when there are heavy conflicts. These kind of answers put your knowledge in question which is not good going into the exam.

upvoted 1 times

pentium75 11 months, 3 weeks ago

The "Correct Answers" for this exam are obviously determined by picking a random letter between A and D.

upvoted 2 times

gx2222 1 year, 8 months ago

Selected Answer: C

To meet the company's requirements of having a resilient solution that can process orders automatically in case of a system outage, the solutions architect needs to implement a fault-tolerant architecture. Based on the given scenario, a potential solution is to move the EC2 instances into an Auto Scaling group and configure the order system to send messages to an Amazon Simple Queue Service (Amazon SQS) queue. The EC2 instances can then consume messages from the queue.

upvoted 3 times

k33 1 year, 8 months ago

Selected Answer: C

Answer : C

upvoted 2 times

nickolaj 1 year, 10 months ago

Selected Answer: C

C. Move the EC2 instances into an Auto Scaling group. Configure the order system to send messages to an Amazon Simple

Queue Service (Amazon SQS) queue. Configure the EC2 instances to consume messages from the queue.

To meet the requirements of the company, a solutions architect should ensure that the system is resilient and can process orders automatically in the event of a system outage. To achieve this, moving the EC2 instances into an Auto Scaling group is a good first step. This will enable the system to automatically add or remove instances based on demand and availability.

upvoted 3 times

nickolaj 1 year, 10 months ago

However, it's also necessary to ensure that orders are not lost if a system outage occurs. To achieve this, the order system can be configured to send messages to an Amazon Simple Queue Service (Amazon SQS) queue. SQS is a highly available and durable messaging service that can help ensure that messages are not lost if the system fails.

Finally, the EC2 instances can be configured to consume messages from the queue, process the orders and then store them in the database on Amazon RDS. This approach ensures that orders are not lost and can be processed automatically if a system outage occurs. Therefore, option C is the correct answer.

upvoted 3 times

nickolaj 1 year, 10 months ago

Option A is incorrect because it suggests creating an Amazon EventBridge rule to target an Amazon Elastic Container Service (ECS) task. While this may be a valid solution in some cases, it is not necessary in this scenario.

Option B is incorrect because it suggests moving the EC2 instances into an Auto Scaling group behind an Application Load Balancer (ALB) and updating the order system to send messages to the ALB endpoint. While this approach can provide resilience and scalability, it does not address the issue of order processing and the need to ensure that orders are not lost if a system outage occurs.

Option D is incorrect because it suggests using Amazon Simple Notification Service (SNS) to send messages to an AWS Lambda function, which will then send a command to the EC2 instances to process the messages by using AWS Systems Manager Run Command. While this approach may work, it is more complex than necessary and does not take advantage of the durability and availability of SQS.

upvoted 3 times

LuckyAro 1 year, 11 months ago

Selected Answer: C

My question is; can orders be sent directly into an SQS queue ? How about the protocol for management of the messages from the queue ? can EC2 instances be programmed to process them like Lambda ?

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #196

Topic 1

A company runs an application on a large fleet of Amazon EC2 instances. The application reads and writes entries into an Amazon DynamoDB table. The size of the DynamoDB table continuously grows, but the application needs only data from the last 30 days. The company needs a solution that minimizes cost and development effort.

Which solution meets these requirements?

- A. Use an AWS CloudFormation template to deploy the complete solution. Redeploy the CloudFormation stack every 30 days, and delete the original stack.
- B. Use an EC2 instance that runs a monitoring application from AWS Marketplace. Configure the monitoring application to use Amazon DynamoDB Streams to store the timestamp when a new item is created in the table. Use a script that runs on the EC2 instance to delete items that have a timestamp that is older than 30 days.
- C. Configure Amazon DynamoDB Streams to invoke an AWS Lambda function when a new item is created in the table. Configure the Lambda function to delete items in the table that are older than 30 days.
- D. Extend the application to add an attribute that has a value of the current timestamp plus 30 days to each new item that is created in the table. Configure DynamoDB to use the attribute as the TTL attribute. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****Gil80** **Highly Voted** 2 years, 6 months ago**Selected Answer: D**

changing my answer to D after researching a bit.

The DynamoDB TTL feature allows you to define a per-item timestamp to determine when an item is no longer needed. Shortly after the date and time of the specified timestamp, DynamoDB deletes the item from your table without consuming any write throughput.

upvoted 42 times

satyaamm **Most Recent** 3 months, 4 weeks ago**Selected Answer: D**

TTL is the most suitable here.

upvoted 1 times

1e22522 10 months, 2 weeks ago

Selected Answer: D

Always bet on the TTL

upvoted 2 times

Nawaff 11 months, 1 week ago

Selected Answer: C

I would day C

Because D requires extending the application to add the timestamp attribute.

Which is by itself a development effort.

upvoted 2 times

Anthony_Rodrigues 9 months, 3 weeks ago

It's to minimize development effort.

C is not the correct one because it needs to enable Dynamo Streams and configure them to send to Lambda, then create the Lambda code for deleting items.

Besides that, the Lambda is triggered only on new items, and since it's ephemeral, it would have to query items older than 30 days and then delete them.

With D, you just need to make two small changes, and it's done—no extra service, no overhead, and no concern about failing to remove.

upvoted 4 times

Hkayne 1 year, 1 month ago

Selected Answer: D

D is the correct answer

upvoted 2 times

soufiyane 1 year, 2 months ago

Selected Answer: D

D is the best answer, dynamostreams is not suitable for this use cases

upvoted 2 times

Uzbekistan 1 year, 2 months ago

Option D is the most suitable solution to meet the company's requirements while minimizing cost and development effort.

TTL (Time to Live) Attribute: DynamoDB provides a feature called Time to Live (TTL), which allows you to automatically delete items from a table after a specified period. By adding a TTL attribute to each item with a value of the current timestamp plus 30 days, you can let DynamoDB automatically delete items older than 30 days. This eliminates the need for manual deletion efforts or periodic stack redeployment.

Minimal Development Effort

Cost-Effective

upvoted 3 times

scar0909 1 year, 3 months ago

Selected Answer: D

use ttl

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

A and B don't solve anything.

Between C and D, C requires more cost due to Lambda executions. D uses the TTL built-in feature so it won't cost extra. Also, D does not require extra development and is a matter of configuration. In old-school developer speak, don't write code if your DBA can do some work!

upvoted 3 times

awsgeek75 1 year, 4 months ago

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/TTL.html>

upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: D

DynamoDB Time to Live was designed to handle this kind of requirement where an item is no longer needed. TTL is provided at no extra cost as a means to reduce stored data volumes by retaining only the items that remain current for your workload's needs

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

D. Extend the application to add an attribute that has a value of the current timestamp plus 30 days to each new item that is created in the table. Configure DynamoDB to use the attribute as the TTL attribute.

The main reasons are:

Using DynamoDB's built-in TTL functionality is the most direct way to handle data expiration.

It avoids the complexity of triggers, streams, and lambda functions in option C.

Modifying the application code to add the TTL attribute is relatively simple and minimizes operational overhead

upvoted 3 times

cookieMr 1 year, 11 months ago

Selected Answer: D

By adding a TTL attribute to the DynamoDB table and setting it to the current timestamp plus 30 days, DynamoDB will automatically delete the items that are older than 30 days. This solution eliminates the need for manual deletion or additional infrastructure components.

A. Redeploying the CloudFormation stack every 30 days and deleting the original stack introduces unnecessary complexity and operational overhead.

B. Using an EC2 instance with a monitoring application and a script to delete items older than 30 days adds additional infrastructure and maintenance efforts.

C. Configuring DynamoDB Streams to invoke a Lambda function to delete items older than 30 days adds complexity and requires additional development and operational effort compared to using the built-in TTL feature of DynamoDB.

upvoted 3 times

pisica134 1 year, 11 months ago

D: <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/TTL.html>

upvoted 2 times

Abrar2022 2 years ago

Amazon DynamoDB Time to Live (TTL) allows you to define a per-item timestamp to determine when an item is no longer needed.

upvoted 4 times

studynoplay 2 years ago

Selected Answer: D

C is incorrect because it can take more than 15 minutes to delete the old data. Lambda won't work

upvoted 2 times

Konb 2 years, 1 month ago

Selected Answer: D

Clear case for TTL - every object gets deleted after a certain period of time

upvoted 2 times

rushi0611 2 years, 1 month ago

Selected Answer: D

Use DynamoDB TTL feature to achieve this..
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #197

Topic 1

A company has a Microsoft .NET application that runs on an on-premises Windows Server. The application stores data by using an Oracle Database Standard Edition server. The company is planning a migration to AWS and wants to minimize development changes while moving the application. The AWS application environment should be highly available.

Which combination of actions should the company take to meet these requirements? (Choose two.)

- A. Refactor the application as serverless with AWS Lambda functions running .NET Core.
- B. Rehost the application in AWS Elastic Beanstalk with the .NET platform in a Multi-AZ deployment. **Most Voted**
- C. Replatform the application to run on Amazon EC2 with the Amazon Linux Amazon Machine Image (AMI).
- D. Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Amazon DynamoDB in a Multi-AZ deployment.
- E. Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Oracle on Amazon RDS in a Multi-AZ deployment. **Most Voted**

Correct Answer: BE*Community vote distribution***Comments****DavidNamy** **Highly Voted** 2 years, 5 months ago**Selected Answer: BE**

B. Rehost the application in AWS Elastic Beanstalk with the .NET platform in a Multi-AZ deployment.
E. Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Oracle on Amazon RDS in a Multi-AZ deployment.

Rehosting the application in Elastic Beanstalk with the .NET platform can minimize development changes. Multi-AZ deployment of Elastic Beanstalk will increase the availability of application, so it meets the requirement of high availability.

Using AWS Database Migration Service (DMS) to migrate the database to Amazon RDS Oracle will ensure compatibility, so the application can continue to use the same database technology, and the development team can use their existing skills. It also migrates to a managed service, which will handle the availability, so the team do not have to worry about it. Multi-AZ deployment will increase the availability of the database.

upvoted 17 times

Buruguduystunstugudunstuy Highly Voted 2 years, 5 months ago

Selected Answer: BE

B. Rehost the application in AWS Elastic Beanstalk with the .NET platform in a Multi-AZ deployment.
E. Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Oracle on Amazon RDS in a Multi-AZ deployment.

To minimize development changes while moving the application to AWS and to ensure a high level of availability, the company can rehost the application in AWS Elastic Beanstalk with the .NET platform in a Multi-AZ deployment. This will allow the application to run in a highly available environment without requiring any changes to the application code.

The company can also use AWS Database Migration Service (AWS DMS) to migrate the Oracle database to Oracle on Amazon RDS in a Multi-AZ deployment. This will allow the company to maintain the existing database platform while still achieving a high level of availability.

upvoted 7 times

lofzee Most Recent 1 year ago

Selected Answer: BE

BE = least effort approach.. basically a lift and shift which is what the questions is asking for

upvoted 2 times

hardy1234567 1 year, 1 month ago

d - incorrect at all. Doesn't exist way for migration oracle to dynamoDB.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: BE

E for minimizing development changes by using same Oracle engine but in highly available deployment.

C and D require platform change so it won't work as it increases development.

A is also development work of converting .Net to .Net core Lambda functions. May not even be possible.

B is simple lift and shift

BE is correct

upvoted 3 times

vijaykamal 1 year, 8 months ago

Selected Answer: BE

DynamoDB is NoSQL - E it out

Replatform requires considerable overhead - C is out

Lambda function is for running code for short duration - A is out

Answer - BE

upvoted 6 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: BE

Minimize development changes + High availability = AWS Elastic Beanstalk and Oracle on Amazon RDS in a Multi-AZ deployment

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

B) Rehost the application in AWS Elastic Beanstalk with the .NET platform in a Multi-AZ deployment.

E) Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Oracle on Amazon RDS in a Multi-AZ deployment.

The reasons are:

° Rehosting in Elastic Beanstalk allows lifting and shifting the .NET application with minimal code changes. Multi-AZ deployment provides high availability.

° Using DMS to migrate the Oracle data to RDS Oracle in Multi-AZ deployment minimizes changes for the database while achieving high availability.

° Together this "lift and shift" approach minimizes refactoring needs while providing HA on AWS.

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: BE

B. This allows the company to migrate the application to AWS without significant code changes while leveraging the scalability and high availability provided by EBS's Multi-AZ deployment.

E. This enables the company to migrate the Oracle database to RDS while maintaining compatibility with the existing application and leveraging the Multi-AZ deployment for high availability.

A. would require significant development changes and may not provide the same level of compatibility as rehosting or replatforming options.

C. would still require changes to the application and the underlying infrastructure, whereas rehosting with EBS minimizes the need for modification.

D. would likely require significant changes to the application code, as DynamoDB is a NoSQL database with a different data model compared to Oracle.

upvoted 4 times

markw92 1 year, 12 months ago

Answer is BE. No idea why D was chosen. That requires development work and question clearly states minimize development changes, changing db from Oracle to DynamoDB is LOT of development.

upvoted 3 times

Bmarodi 2 years ago

Selected Answer: BE

B + E are the answers that fulfil the requirements.

upvoted 2 times

cheese929 2 years, 1 month ago

Selected Answer: BE

B and E

upvoted 2 times

Nikhilcy 2 years, 1 month ago

why not C?

upvoted 2 times

AWSSURI 9 months, 2 weeks ago

Retire(simplest) < Retain < Relocate < Rehost < Repurchase < Replatform < Re-architect/Refactor (most complex)....So by this order rehosting less complex than replatform thats why we go for BE

upvoted 2 times

AlankarJ 2 years ago

It runs on a windows server, shifting the whole this to Linux based EC2 would be extra work and of no sense

upvoted 3 times

k33 2 years, 2 months ago

Selected Answer: BE

Answer : BE

upvoted 2 times

waiyu9981 2 years, 5 months ago

Why A is wrong?

upvoted 1 times

gustavtd 2 years, 5 months ago

Because that needs some development,

upvoted 3 times

upvoted 2 times

techhb 2 years, 5 months ago

Selected Answer: BE

B&E Option ,because D is for No-Sql
upvoted 2 times

JayBee65 2 years, 5 months ago

And requires additional development effort
upvoted 2 times

career360guru 2 years, 5 months ago

B&E Option
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #198

Topic 1

A company runs a containerized application on a Kubernetes cluster in an on-premises data center. The company is using a MongoDB database for data storage. The company wants to migrate some of these environments to AWS, but no code changes or deployment method changes are possible at this time. The company needs a solution that minimizes operational overhead.

Which solution meets these requirements?

- A. Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 worker nodes for compute and MongoDB on EC2 for data storage.
- B. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate for compute and Amazon DynamoDB for data storage
- C. Use Amazon Elastic Kubernetes Service (Amazon EKS) with Amazon EC2 worker nodes for compute and Amazon DynamoDB for data storage.
- D. Use Amazon Elastic Kubernetes Service (Amazon EKS) with AWS Fargate for compute and Amazon DocumentDB (with MongoDB compatibility) for data storage. **Most Voted**

Correct Answer: D*Community vote distribution*

D (100%)

Comments**Marge_Simpson** **Highly Voted** 2 years ago**Selected Answer: D**

If you see MongoDB, just go ahead and look for the answer that says DocumentDB.
upvoted 40 times

Guru4Cloud **Highly Voted** 1 year, 3 months ago**Selected Answer: D**

Option D is the correct solution that meets all the requirements:

- ° Use Amazon Elastic Kubernetes Service (Amazon EKS) with AWS Fargate for compute and Amazon DocumentDB (with MongoDB compatibility) for data storage.

The key reasons are:

- ° EKS allows running the Kubernetes environment on AWS without changes.

- Using Fargate removes the need to provision and manage EC2 instances.
 - DocumentDB provides MongoDB compatibility so the data layer is unchanged.
- upvoted 10 times

MehulKapadia Most Recent 8 months ago

Selected Answer: D

Applications are already containerized. Amazon EKS is fully managed kubernetes service.
FarGate = Less overhead of managing infrastructure.
Amazon DocumentDB is MongoDB Compatible.

Answer D

upvoted 2 times

LoXoL 11 months ago

Selected Answer: D

no brainer says D

upvoted 2 times

james2033 1 year, 4 months ago

Selected Answer: D

Question keyword "containerized application", "Kubernetes cluster", "no changes or deployment method changes". Choose C, not D.

But "minimizes operational overhead", choose D.

upvoted 2 times

cookieMr 1 year, 5 months ago

Selected Answer: D

This solution allows the company to leverage EKS to manage the K8s cluster and Fargate to handle the compute resources without requiring manual management of EC2 worker nodes. The use of DocumentDB provides a fully managed MongoDB-compatible database service in AWS.

A. would require managing and scaling the EC2 instances manually, which increases operational overhead.

B. would require significant changes to the application code as DynamoDB is a NoSQL database with a different data model compared to MongoDB.

C. would also require code changes to adapt to DynamoDB's different data model, and managing EC2 worker nodes increases operational overhead.

upvoted 4 times

Bmarodi 1 year, 6 months ago

Selected Answer: D

The solution meets these requirements is option D.

upvoted 2 times

studynoplay 1 year, 7 months ago

Selected Answer: D

minimizes operational overhead = Serverless (Fargate)

MongoDB = DocumentDB

upvoted 2 times

Buruguduystunstugudunstuy 1 year, 11 months ago

Selected Answer: D

To minimize operational overhead and avoid making any code or deployment method changes, the company can use Amazon Elastic Kubernetes Service (EKS) with AWS Fargate for computing and Amazon DocumentDB (with MongoDB compatibility) for data storage. This solution allows the company to run the containerized application on EKS without having to manage the underlying infrastructure or make any changes to the application code.

AWS Fargate is a fully-managed container execution environment that allows you to run containerized applications without the need to manage the underlying EC2 instances.

Amazon DocumentDB is a fully-managed document database service that supports MongoDB workloads, allowing the company to use the same database platform as in their on-premises environment without having to make any code changes.
upvoted 5 times

techhb 1 year, 11 months ago

Selected Answer: D

Reason A & B Eliminated as its Kubernetes
why D read here <https://containersonaws.com/introduction/ec2-or-aws-fargate/>
upvoted 3 times

career360guru 1 year, 12 months ago

Selected Answer: D

Option D
upvoted 3 times

dcyberguy 2 years ago

DDDDDDD
upvoted 2 times

Gabs90 2 years ago

Selected Answer: D

<https://www.examttopics.com/discussions/amazon/view/67897-exam-aws-certified-solutions-architect-associate-saa-c02/>
upvoted 2 times

leonnann 2 years ago

Selected Answer: D

D meets the requirements
upvoted 2 times

Nigma 2 years ago

Selected Answer: D

D
EKS because of Kubernetes so A and B are eliminated
not C because of MongoDB and Fargate is more expensive
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #199

Topic 1

A telemarketing company is designing its customer call center functionality on AWS. The company needs a solution that provides multiple speaker recognition and generates transcript files. The company wants to query the transcript files to analyze the business patterns. The transcript files must be stored for 7 years for auditing purposes.

Which solution will meet these requirements?

- A. Use Amazon Rekognition for multiple speaker recognition. Store the transcript files in Amazon S3. Use machine learning models for transcript file analysis.
- B. Use Amazon Transcribe for multiple speaker recognition. Use Amazon Athena for transcript file analysis. **Most Voted**
- C. Use Amazon Translate for multiple speaker recognition. Store the transcript files in Amazon Redshift. Use SQL queries for transcript file analysis.
- D. Use Amazon Rekognition for multiple speaker recognition. Store the transcript files in Amazon S3. Use Amazon Textract for transcript file analysis.

Correct Answer: B*Community vote distribution*

B (93%)

Other (7%)

Comments**Buruguduystunstugudunstuy** **Highly Voted** 1 year, 11 months ago**Selected Answer: B**

The correct answer is B: Use Amazon Transcribe for multiple speaker recognition. Use Amazon Athena for transcript file analysis.

Amazon Transcribe is a service that automatically transcribes spoken language into written text. It can handle multiple speakers and can generate transcript files in real-time or asynchronously. These transcript files can be stored in Amazon S3 for long-term storage.

Amazon Athena is a query service that allows you to analyze data stored in Amazon S3 using SQL. You can use it to analyze the transcript files and identify patterns in the data.

Option A is incorrect because Amazon Rekognition is a service for analyzing images and videos, not transcribing spoken language.

Option C is incorrect because Amazon Translate is a service for translating text from one language to another, not transcribing spoken language.

Option D is incorrect because Amazon Textract is a service for extracting text and data from documents and images, not transcribing spoken language.

upvoted 26 times

TheAbsoluteTruth 1 year, 8 months ago

What bothers me is the 7 years of storage.

upvoted 8 times

enzomv 1 year, 10 months ago

The correct answer is C.

<https://docs.aws.amazon.com/transcribe/latest/dg/what-is.html>

You can transcribe streaming media in real time or you can upload and transcribe media files. To see which languages are supported for each type of transcription, refer to the Supported languages and language-specific features table.

upvoted 4 times

enzomv 1 year, 10 months ago

Disregard. I meant B

upvoted 2 times

enzomv 1 year, 10 months ago

<https://aws.amazon.com/about-aws/whats-new/2022/06/amazon-transcribe-supports-automatic-language-identification-multi-lingual-audio/>

Amazon Transcribe is a service for multi-language identification, which identifies all languages spoken in the audio file and creates transcript using each identified language.

upvoted 2 times

enzomv 1 year, 10 months ago

Disregard. I meant Amazon Transcribe

upvoted 1 times

awsgeek75 Highly Voted 11 months, 2 weeks ago

Selected Answer: B

This is a poorly worded question with poorly worded options. Rekognition and Translate cannot convert speech to text so those options A, C & D are gone. B is the closes option but it does not mention S3 or retention policy of 7 years. Just a best guess on massive assumptions.

upvoted 7 times

bora4motion Most Recent 2 weeks, 5 days ago

Selected Answer: B

B - amazon transcribe is the right tool for call centers

upvoted 1 times

satyaamm 3 months, 4 weeks ago

Selected Answer: B

Transcribe for speech to text conversion and Athena for analysis.

upvoted 1 times

james_3005 8 months, 2 weeks ago

Option B is the most relevant one, but it doesn't mention how to retain data in 7 years...

upvoted 2 times

SinghJagdeep 11 months, 2 weeks ago

Selected Answer: B

check out this blog here: <https://aws.amazon.com/de/blogs/machine-learning/automating-the-analysis-of-multi-speaker-audio-files-using-amazon-transcribe-and-amazon-athena/>

upvoted 2 times

upvoted 3 times

pentium75 11 months, 3 weeks ago

Selected Answer: B

Perfectly explained here: <https://aws.amazon.com/de/blogs/machine-learning/automating-the-analysis-of-multi-speaker-audio-files-using-amazon-transcribe-and-amazon-athena/>

upvoted 4 times

yodelin 1 year, 2 months ago

really hope I could have this kind of question during the exam, 4 different techs in the first 5 words of the answer! Just go with the correct one and ignore the rest of the text XDDD

upvoted 3 times

paniya93 1 year, 2 months ago

Selected Answer: B

<https://aws.amazon.com/blogs/machine-learning/automating-the-analysis-of-multi-speaker-audio-files-using-amazon-transcribe-and-amazon-athena/>

upvoted 2 times

vijaykamal 1 year, 2 months ago

Selected Answer: B

Amazon Rekognition is primarily designed for image and video analysis, not for transcribing audio or recognizing multiple speakers. -> Option A and D are ruled out

Amazon Translate is used for language translation -> Option C is ruled out

upvoted 3 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: B

Provide multiple speaker recognition and generate transcript files = Amazon Transcribe

Query the transcript files = Amazon Athena

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

The correct answer is B: Use Amazon Transcribe for multiple speaker recognition. Use Amazon Athena for transcript file analysis.

upvoted 3 times

Thornessen 1 year, 4 months ago

Selected Answer: B

Tricky or incomplete question..

B is the answer because Transcribe is the right service for processing voice calls.

But 7 years of storage is not covered (should add S3 storage)

And Athena querying is just SQL querying, it cannot help you much to recognize business patterns, for that I would think some text analysis service like Comprehend would be needed.

Unless... We use Transcribe not only to transcribe, but also to recognize some key words, and then create a DB/S3 record with multiple fields, e.g. if it is a telemarketing questionnaire, record answer for each question. Then SQL querying might be useful.

upvoted 3 times

pentium75 11 months, 3 weeks ago

"The company wants to query the transcript files" is the requirement. How they will be using the query results "to analyze the business patterns" is not our issue.

The "7 years" are not mentioned in any of the options, but Transcribe stores results in S3.

upvoted 2 times

sickcow 1 year, 5 months ago

Selected Answer: C

Transcribe and (s3) + Athena is the way to go here.
Redshift sounds like an overkill

upvoted 3 times

cookieMr 1 year, 5 months ago

Amazon Transcribe provides accurate transcription of audio recordings with multiple speakers, generating transcript files. These files can be stored in Amazon S3. To analyze the transcripts and extract insights, Amazon Athena allows SQL-based querying of the stored files.

A. Amazon Rekognition is for image and video analysis, not audio transcription.

C. Amazon Translate is for language translation, not speaker recognition or transcript analysis. Amazon Redshift may not be the best choice for storing and querying transcript files.

D. Amazon Rekognition is for image and video analysis, and Amazon Textract is for document extraction, not suitable for audio transcription or analysis. Storing the transcript files in S3 is appropriate, but the analysis requires a different service like Amazon Athena.

upvoted 2 times

Bmarodi 1 year, 6 months ago

Selected Answer: B

the solution that meets these requirements is option B.

upvoted 2 times

cheese929 1 year, 7 months ago

Selected Answer: B

B is correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #203

Topic 1

The customers of a finance company request appointments with financial advisors by sending text messages. A web application that runs on Amazon EC2 instances accepts the appointment requests. The text messages are published to an Amazon Simple Queue Service (Amazon SQS) queue through the web application. Another application that runs on EC2 instances then sends meeting invitations and meeting confirmation email messages to the customers. After successful scheduling, this application stores the meeting information in an Amazon DynamoDB database.

As the company expands, customers report that their meeting invitations are taking longer to arrive.

What should a solutions architect recommend to resolve this issue?

- A. Add a DynamoDB Accelerator (DAX) cluster in front of the DynamoDB database.
- B. Add an Amazon API Gateway API in front of the web application that accepts the appointment requests.
- C. Add an Amazon CloudFront distribution. Set the origin as the web application that accepts the appointment requests.
- D. Add an Auto Scaling group for the application that sends meeting invitations. Configure the Auto Scaling group to scale based on the depth of the SQS queue. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Buruguduystunstugudunstuy **Highly Voted** 1 year, 11 months ago

Selected Answer: D

Option D. Add an Auto Scaling group for the application that sends meeting invitations. Configure the Auto Scaling group to scale based on the depth of the SQS queue.

To resolve the issue of longer delivery times for meeting invitations, the solutions architect can recommend adding an Auto Scaling group for the application that sends meeting invitations and configuring the Auto Scaling group to scale based on the depth of the SQS queue. This will allow the application to scale up as the number of appointment requests increases, improving the performance and delivery times of the meeting invitations.

upvoted 17 times

cookieMr **Highly Voted** 1 year, 5 months ago

Selected Answer: D

By adding an ASG for the application that sends meeting invitations and configuring it to scale based on the depth of the SQS, the system can automatically adjust its capacity based on the number of pending messages in the queue. This ensures that the application can handle increased message load and process the meeting invitations more efficiently, reducing the delay experienced by customers.

A. Adding a DynamoDB Accelerator (DAX) cluster in front of the DynamoDB database would improve read performance for DynamoDB, but it does not directly address the issue of delayed meeting invitations.

B. Adding an API Gateway API in front of the web application that accepts the appointment requests may help with request handling and management, but it does not directly address the issue of delayed meeting invitations.

C. Adding a CloudFront distribution with the web application as the origin would improve content delivery and caching, but it does not directly address the issue of delayed meeting invitations.

upvoted 11 times

67a3f49 **Most Recent** 9 months, 3 weeks ago

First question with consistent answer :)

upvoted 5 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: D

Add an Auto Scaling group for the application that sends meeting invitations. Configure the Auto Scaling group to scale based on the depth of the SQS queue.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

Add an Auto Scaling group for the application that sends meeting invitations. Configure the Auto Scaling group to scale based on the depth of the SQS queue.

upvoted 2 times

career360guru 1 year, 12 months ago

Selected Answer: D

Option D is the right Answer,

upvoted 3 times

k1kavi1 1 year, 12 months ago

Selected Answer: D

Agreed

upvoted 2 times

jambajuice 2 years ago

Selected Answer: D

ANswer d

upvoted 2 times

leonnnn 2 years ago

Selected Answer: D

D meets the requirements

upvoted 2 times

Nigma 2 years ago

Selected Answer: D

Answer : D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #204

Topic 1

An online retail company has more than 50 million active customers and receives more than 25,000 orders each day. The company collects purchase data for customers and stores this data in Amazon S3. Additional customer data is stored in Amazon RDS.

The company wants to make all the data available to various teams so that the teams can perform analytics. The solution must provide the ability to manage fine-grained permissions for the data and must minimize operational overhead.

Which solution will meet these requirements?

- A. Migrate the purchase data to write directly to Amazon RDS. Use RDS access controls to limit access.
- B. Schedule an AWS Lambda function to periodically copy data from Amazon RDS to Amazon S3. Create an AWS Glue crawler. Use Amazon Athena to query the data. Use S3 policies to limit access.
- C. Create a data lake by using AWS Lake Formation. Create an AWS Glue JDBC connection to Amazon RDS. Register the S3 bucket in Lake Formation. Use Lake Formation access controls to limit access. **Most Voted**
- D. Create an Amazon Redshift cluster. Schedule an AWS Lambda function to periodically copy data from Amazon S3 and Amazon RDS to Amazon Redshift. Use Amazon Redshift access controls to limit access.

Correct Answer: C

Community vote distribution

C (100%)

Comments

anhike **Highly Voted** 2 years, 5 months ago

Answer : C keyword "manage-fine-grained"

<https://aws.amazon.com/blogs/big-data/manage-fine-grained-access-control-using-aws-lake-formation/>

upvoted 27 times

markw92 1 year, 11 months ago

You can manage fine grained using redshift as well - <https://aws.amazon.com/blogs/big-data/achieve-fine-grained-data-security-with-row-level-access-control-in-amazon-redshift/>

But, I believe the keyword to look for is "minimize operational overhead", which lakeformation does without duplicating much of the data. Redshift is operational overhead and duplication of data. not sure whv the answer is D. i vote C as well.

of the data requires operational overhead and duplication of data between any one channel to another can result in data being upvoted 9 times

Olaunfazed 1 year, 11 months ago

yeah, most of examtopics answers are wrong
upvoted 9 times

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: C

Option C. Create a data lake by using AWS Lake Formation. Create an AWS Glue JDBC connection to Amazon RDS. Register the S3 bucket in Lake Formation. Use Lake Formation access controls to limit access.

To make all the data available to various teams and minimize operational overhead, the company can create a data lake by using AWS Lake Formation. This will allow the company to centralize all the data in one place and use fine-grained access controls to manage access to the data.

To meet the requirements of the company, the solutions architect can create a data lake by using AWS Lake Formation, create an AWS Glue JDBC connection to Amazon RDS, and register the S3 bucket in Lake Formation. The solutions architect can then use Lake Formation access controls to limit access to the data. This solution will provide the ability to manage fine-grained permissions for the data and minimize operational overhead.

upvoted 5 times

majdango 2 years ago

.....

upvoted 1 times

ike001 **Most Recent** 11 months, 3 weeks ago

Answer is C
upvoted 3 times

MehulKapadia 1 year, 1 month ago

Selected Answer: C

Keyword: "manage fine-grained permissions for data"
Data Lake Using Lake Formation: manage fine-grained permissions for the data with ease.

Fine grained permissions for data = Lake Formation

Answer: C
upvoted 4 times

LoXoL 1 year, 4 months ago

Selected Answer: C

C represents the easiest way to ingest data from S3 and control accesses.
upvoted 1 times

karloscetina007 1 year, 8 months ago

Selected Answer: C

a fine grained permissons is one of the conditions to accomplishes with the requirement.
With the use of AWS Glue you can get accomplishes with this requirement.
My answer is: C
upvoted 4 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: C

With Lake formation you can scale permissions more easily with fine-grained security capabilities, including row- and cell-level permissions and tag-based access control.
upvoted 5 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: C

Lake Formation enables the creation of a secure and scalable data lake on AWS, allowing centralized access controls for both S3 and RDS data. By using Lake Formation, the company can manage permissions effectively and integrate RDS data through the AWS Glue JDBC connection. Registering the S3 in Lake Formation ensures unified access control. This solution reduces operational overhead while providing fine-grained permissions management.

upvoted 4 times

cookieMr 1 year, 11 months ago

Selected Answer: C

Lake Formation enables the creation of a secure and scalable data lake on AWS, allowing centralized access controls for both S3 and RDS data. By using Lake Formation, the company can manage permissions effectively and integrate RDS data through the AWS Glue JDBC connection. Registering the S3 in Lake Formation ensures unified access control. This solution reduces operational overhead while providing fine-grained permissions management.

A. Directly writing purchase data to Amazon RDS with RDS access controls lacks comprehensive permissions management for both S3 and RDS data.

B. Periodically copying data from RDS to S3 using Lambda and using AWS Glue and Athena for querying does not offer fine-grained permissions management and introduces data synchronization complexities.

D. Creating an Redshift cluster and copying data from S3 and RDS to Redshift adds complexity and operational overhead without the flexibility of Lake Formation's permissions management capabilities.

upvoted 4 times

pisica134 1 year, 11 months ago

Answer is C AWS Lake Formation provides a comprehensive solution for building and managing a data lake. It simplifies data ingestion, organization, and access control. By creating a data lake using AWS Lake Formation, you can centralize and govern access to your data across multiple sources.

upvoted 2 times

Bmarodi 1 year, 12 months ago

Selected Answer: C

Option C is right answer: <https://docs.aws.amazon.com/lake-formation/latest/dg/what-is-lake-formation.html>

upvoted 2 times

Abrar2022 2 years ago

Lake Formation helps you manage fine-grained access for internal and external customers from a centralized location and in a scalable way.

upvoted 2 times

doorahmie 2 years, 4 months ago

<https://docs.aws.amazon.com/lake-formation/latest/dg/access-control-overview.html>

upvoted 3 times

LuckyAro 2 years, 4 months ago

Selected Answer: C

To me, the give-away was: "The company wants to make all the data available to various teams" - Data-Lake - All data in one place.

upvoted 5 times

master1004 2 years, 5 months ago

The correct answer is D.

The company uses all the data from various teams so that the teams can do their analysis.

Therefore, it is the best way to separately configure redshift for data warehousing and for all employees to connect to the redshift DB and perform analysis tasks without burdening the operating DB (must minimize operational overhead).

upvoted 4 times

ruqui 2 years ago

I don't think that "periodically copy data from Amazon S3 and RDS to Redshift" minimize the operational overhead. The correct answer for me is C

upvoted 2 times

aba2s 2 years, 5 months ago

Selected Answer: C

Manage fine-grained access control using AWS Lake Formation

<https://aws.amazon.com/blogs/big-data/manage-fine-grained-access-control-using-aws-lake-formation/>

upvoted 2 times

kvenikoduru 2 years, 5 months ago

Selected Answer: C

a combination of the following 2 URLs I believe it is C

<https://aws.amazon.com/lake-formation/>

<https://aws.amazon.com/blogs/big-data/manage-fine-grained-access-control-using-aws-lake-formation/>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #205


Topic 1

A company hosts a marketing website in an on-premises data center. The website consists of static documents and runs on a single server. An administrator updates the website content infrequently and uses an SFTP client to upload new documents.

The company decides to host its website on AWS and to use Amazon CloudFront. The company's solutions architect creates a CloudFront distribution. The solutions architect must design the most cost-effective and resilient architecture for website hosting to serve as the CloudFront origin.

Which solution will meet these requirements?

- A. Create a virtual server by using Amazon Lightsail. Configure the web server in the Lightsail instance. Upload website content by using an SFTP client.
- B. Create an AWS Auto Scaling group for Amazon EC2 instances. Use an Application Load Balancer. Upload website content by using an SFTP client.
- C. Create a private Amazon S3 bucket. Use an S3 bucket policy to allow access from a CloudFront origin access identity (OAI). Upload website content by using the AWS CLI. **Most Voted**
- D. Create a public Amazon S3 bucket. Configure AWS Transfer for SFTP. Configure the S3 bucket for website hosting. Upload website content by using the SFTP client.

Correct Answer: C*Community vote distribution***Comments****bjexamprep** **Highly Voted** 1 year, 10 months ago**Selected Answer: C**

The question here is whether the solution architect can change the requirement. The requirement says very clear about SFTP which cannot be addressed by option C. But the question also gives very clear hint about OAI which cannot be addressed by option D. Option D also doesn't mention anything about CloudFront which is part of the requirement of the question. So, if the requirement cannot be changed, D is the answer; if the requirement can be changed, C is the answer. But if the requirement can be changed, what's the limitation? That will be a Chaos. I'm voting C, and curse the question designer.
upvoted 24 times

Iconique 1 year, 8 months ago

"The solutions architect must design the most cost-effective and resilient architecture for website hosting to serve as the CloudFront origin." The solution architect is looking for a solution that can fit with CloudFront as origin! So it doesn't matter that option D does not mention CF, CF is part of the solution!
Having a marketing website on-premise clearly indicates having S3 as static content.
AWS Transfer Family is the way to upload files via FTP to S3!
So the answer is D.

Why not C?

User is already uploading content via FTP, option C is eliminating this option for him and forces using the CLI. The solution from C does not meet the requirements of having FTP.

upvoted 8 times

Tsige 7 months, 2 weeks ago

The main reason Option D is not the best answer is because it involves additional complexity and cost compared to Option C. Here's a breakdown:

Option D requires setting up AWS Transfer for SFTP, which adds an extra service to manage and incurs additional costs. While it allows the administrator to continue using the SFTP client, it introduces more operational overhead.

Option C leverages Amazon S3 and CloudFront directly, which is simpler and more cost-effective. By using an S3 bucket policy and CloudFront origin access identity (OAI), you ensure secure access without the need for additional services.

In summary, Option C provides a more streamlined and cost-effective solution for hosting the website and serving as the CloudFront origin.

upvoted 4 times

cookieMr **Highly Voted** 1 year, 11 months ago

Selected Answer: C

Hosting the website in a private S3 provides cost-effective and highly available storage for the static website content. By configuring a bucket policy to allow access from a CloudFront OAI, the S3 can be securely accessed only through CloudFront. This ensures that the website content is served through CloudFront while keeping the S3 private. Uploading website content using the AWS CLI allows for easy and efficient content management.

A. Hosting the website on an Lightsail virtual server would introduce additional management overhead and costs compared to using S3 directly for static content hosting.

B. Using an AWS ASG with EC2 instances and an ALB is not necessary for serving static website content. It would add unnecessary complexity and cost.

D. While using AWS Transfer for SFTP allows for SFTP uploads, it introduces additional costs and complexity compared to directly uploading content to an S3 using the AWS CLI. Additionally, hosting the website content in a public S3 may not be desirable from a security standpoint.

upvoted 9 times

jayessh **Most Recent** 6 months, 3 weeks ago

Selected Answer: C

Options A and B involve more complexity and higher costs due to the need for managing instances and load balancers. Option D, while feasible, involves making the S3 bucket public, which is less secure compared to using a private bucket with an OAI.

correct Answer : C

upvoted 2 times

JohnYu 8 months ago

the option D creating a public S3 bucket could expose the website content to the public, which is a security risk. Using an OAI with a private S3 bucket (as in Option C) is more secure. AWS Transfer for SFTP adds unnecessary complexity unless SFTP specifically is required.

upvoted 2 times

MatAlves 9 months ago

Selected Answer: C

"The company... use Amazon CloudFront" = C is the only option that mentions CloudFront.

upvoted 2 times

Pr1est 1 year, 1 month ago

Selected Answer: C

This is another great example of how AWS creates crappy tests. Even internal Tests for employees have so many flaws that people is always creating tickets challenging Questions poorly worded.

upvoted 5 times

f761d0e 1 year, 1 month ago

Another reason why A is better than C:

" OAC helps you secure your origins, such as for Amazon S3. We recommend using OAC"

" If your origin is an Amazon S3 bucket configured as a website endpoint, you must set it up with CloudFront as a custom origin. That means you can't use OAC (or OAI)"

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/private-content-restricting-access-to-s3.html>

upvoted 1 times

djgodzilla 1 year, 5 months ago

Selected Answer: D

you can see in this figure that transfer family framework allows for the data to be available for a broad variety of use cases including content distribution (CF)

<https://d1.awsstatic.com/HiW%20SFTP%20Connectors%20v3.920176622d281d0bd087518827314169b496a055.png>

upvoted 2 times

awsgeek75 1 year, 4 months ago

Two main problems with D:

It's public S3 behind CloudFront which is not well-architected

Infrequent site updates using SFTP so with S3 it will be cli changes. They don't need fancy transfer for this. Right?

upvoted 2 times

MiniYang 1 year, 6 months ago

Selected Answer: C

I think the this is a big misleading " SFTP" (doesn't usally upload) , and it said clearly need Cloudfront and want a cheep solution. So I chose "C"

upvoted 3 times

rlamberti 1 year, 7 months ago

Selected Answer: C

Transferring via AWS CLI is cheaper than via Transfer Family.

It is not the best option, but will do the job of uploading the data to S3.

upvoted 2 times

juanrasus2 1 year, 7 months ago

I'd go with D. In C there is no mention to S3 bucket being configured for web hosting. Simply adding the Cloudfront distribution and pointing that to the S3 won't work out of the box.

upvoted 2 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: D

D - SFTP client to upload new documents.

upvoted 2 times

baku98 1 year, 5 months ago

D is the only one possible.

C cannot be because: In Amazon CloudFront: For Restricting access to an Amazon S3 origin: If your origin is an Amazon S3 bucket configured as a website endpoint, you must set it up with CloudFront as a custom origin. That means you can't use OAC (or OAI).

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/private-content-restricting-access-to-s3.html>

upvoted 2 times

Guru4Cloud 1 year, 8 months ago

I changed C. is better then D

upvoted 4 times

upvoted 4 times

eugene_stalker 2 years ago

Selected Answer: D

D - SFTP client to upload new documents.

upvoted 2 times

bdp123 2 years, 3 months ago

Selected Answer: C

AWS transfer is a cost and doesn't mention using CloudFront

<https://aws.amazon.com/aws-transfer-family/pricing/>

upvoted 4 times

Yelizaveta 2 years, 3 months ago

Selected Answer: C

If you don't want to disable block public access settings for your bucket but you still want your website to be public, you can create a Amazon CloudFront distribution to serve your static website. For more information, see Use an Amazon CloudFront distribution to serve a static website in the Amazon Route 53 Developer Guide.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteAccessPermissionsReqd.html>

upvoted 3 times

PDR 2 years, 4 months ago

Selected Answer: C

I at first thought D but it is in fact C because

"D: Create a public Amazon S3 bucket. Configure AWS Transfer for SFTP. Configure the S3 bucket for website hosting. Upload website content by using the SFTP client." questions says that the company has decided to use Amazon Cloudfront and this answer does not reference using CF and setting S3 as the Origin

"C. Create a private Amazon S3 bucket. Use an S3 bucket policy to allow access from a CloudFront origin access identity (OAI). Upload website content by using the AWS CLI." - mentions CF and the origin and the AWS CLI does infact support transfer by SFTP (which was the part I originally doubted but this link evidences that it does:

<https://docs.aws.amazon.com/cli/latest/reference/transfer/describe-server.html>

upvoted 5 times

bullrem 2 years, 4 months ago

Selected Answer: D

Option C, creating a private Amazon S3 bucket and using an S3 bucket policy to allow access from a CloudFront origin access identity (OAI), would not be the most cost-effective solution. While it would allow the company to use Amazon S3 for storage, it would also require additional setup and maintenance of the OAI, which would add additional cost. Additionally, this solution would not allow the use of SFTP client for uploading content which is the current method used by the company.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #206

Topic 1

A company wants to manage Amazon Machine Images (AMIs). The company currently copies AMIs to the same AWS Region where the AMIs were created. The company needs to design an application that captures AWS API calls and sends alerts whenever the Amazon EC2 CreateImage API operation is called within the company's account.


Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an AWS Lambda function to query AWS CloudTrail logs and to send an alert when a CreateImage API call is detected.
- B. Configure AWS CloudTrail with an Amazon Simple Notification Service (Amazon SNS) notification that occurs when updated logs are sent to Amazon S3. Use Amazon Athena to create a new table and to query on CreateImage when an API call is detected.

C. Create an Amazon EventBridge (Amazon CloudWatch Events) rule for the CreateImage API call. Configure the target as an Amazon Simple Notification Service (Amazon SNS) topic to send an alert when a CreateImage API call is detected.

Most Voted

- D. Configure an Amazon Simple Queue Service (Amazon SQS) FIFO queue as a target for AWS CloudTrail logs. Create an AWS Lambda function to send an alert to an Amazon Simple Notification Service (Amazon SNS) topic when a CreateImage API call is detected.

Correct Answer: C*Community vote distribution***Comments****[Removed]** **Highly Voted** 2 years, 6 months ago**Selected Answer: C**

I'm team C.

<https://docs.aws.amazon.com/AWSEC2/latest/WindowsGuide/monitor-ami-events.html#:~:text=For%20example%2C%20you%20can%20create%20an%20EventBridge%20rule%20that%20detects%20when%20the%20AMI%20creation%20process%20has%20completed%20and%20then%20invokes%20an%20Amazon%20SNS%20topic%20to%20send%20an%20email%20notification%20to%20you.>

upvoted 23 times

MutiverseAgent 1 year, 10 months ago

C is correct > <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/monitor-ami-events.html>
upvoted 2 times

JayBee65 2 years, 4 months ago

That link contains the exact use case and explains how C can be used.
Option B requires you to send logs to S3 and use Athena, 2 additional services that are not required, so this does not meet the "LEAST operational overhead?" requirement, since these are extra services requiring management.
upvoted 6 times

Wajif Highly Voted 2 years, 5 months ago

Selected Answer: A

Why not A? API calls are already logged in Cloudtrail.
upvoted 15 times

pentium75 1 year, 5 months ago

"Least operational overhead" is when the event triggers an action, not when you run a scheduled task that searches logs for the event.
upvoted 6 times

lofzee Most Recent 1 year ago

Selected Answer: C

C all day. Trust me
upvoted 3 times

ManikRoy 1 year, 1 month ago

Selected Answer: C

It can be done with option A but you'll have to write a lambda function. option C is least operational overhead.
upvoted 3 times

zinabu 1 year, 1 month ago

Selected Answer: C

Monitor AMI events using Amazon EventBridge is possible and here is the link:
<https://docs.aws.amazon.com/AWSEC2/latest/WindowsGuide/monitor-ami-events.html>
upvoted 2 times

Mahmoudddddddd 1 year, 2 months ago

Just took the exam today, most of the questions were from here wish I saw them all to be honest before entering the exam. Anyways, this question was at the exam, I picked option A because as the question stated it wanted two things not one thing only an application that CAPTURES API calls and SEND ALERTS WHENEVER CreateImage API call is made, OPTION C CLEARLY STATES THAT IN THIS CASE IT WILL ONLYYY LOOK FOR CREATEIMAGE API CALL it will not capture other API calls like the lambda in option A would! Am i the only one that is thinks that or what? TBH I am not sure about anything in this question but that is why I did not pick option C during the exam.
upvoted 3 times

JA2018 6 months, 3 weeks ago

Hi Mahmoudddddddd, care to share which answer did you actually choose in your exam? Thanks.
upvoted 1 times

bujuman 1 year, 4 months ago

Selected Answer: D

On of the requirements is LEAST operational overhead
CloudTrail sends a notification when log files are written to the Amazon S3 bucket. An active account can generate a large number of notifications. If you subscribe with email or SMS, you can receive a large volume of messages. We recommend that you subscribe using Amazon Simple Queue Service (Amazon SQS), which lets you handle notifications programmatically. For more information, see Subscribing a Queue to an Amazon SNS Topic in the Amazon Simple Queue Service Developer Guide.
upvoted 1 times

Wang87 1 year, 5 months ago

Selected Answer: C

Answer is c.

upvoted 2 times

farnamjam 1 year, 5 months ago

Selected Answer: C

C is correct

upvoted 2 times

master9 1 year, 5 months ago

Selected Answer: D

AWS CloudTrail primarily focuses on auditing and recording API calls made in your AWS account. It logs all API requests made via the AWS Management Console, AWS SDKs, command line tools, and other AWS services. This includes the identity of the caller, the time of the API call, the source IP address of the caller, the request parameters, and the response elements returned by the AWS service. This information is useful for security analysis, resource change tracking, and troubleshooting.

upvoted 1 times

pentium75 1 year, 5 months ago

But this is not about "auditing and recording", you don't want to create reports who created images during the last year, you want an instant alert when someone creates an image. Thus CloudWatch Events.

upvoted 2 times

[Removed] 1 year, 5 months ago

Cloud Watch = AWS Monitoring service for any AWS resources

Cloud Trail = AWS API monitoring service with respect to application event that are hosted on AWS.

Answer would be "C"

<https://docs.aws.amazon.com/AWSEC2/latest/WindowsGuide/monitor-ami-events.html>
service

upvoted 5 times

rlamberti 1 year, 7 months ago

Selected Answer: C

"LEAST operational overhead"

Option A involves coding a Lambda. Not good!

Option C seems to be the correct.

upvoted 3 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: C

Event bridge was built specifically to handle this kind of scenario:

CreateImage API call (Event Source) -> Event bus -> Rules -> Amazon SNS (Event target)

upvoted 7 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

C. Create an Amazon EventBridge (Amazon CloudWatch Events) rule for the CreateImage API call. Configure the target as an Amazon Simple Notification Service (Amazon SNS) topic to send an alert when a CreateImage API call is detected

upvoted 4 times

Nava702 1 year, 9 months ago

Selected Answer: A

A look like the least overhead option to capture an API call.

upvoted 2 times

Mia2009687 1 year, 11 months ago

Selected Answer: B

The company needs to design an application that captures AWS API calls and sends alerts whenever the Amazon EC2 CreateImage API operation is called within the company's account.

With option C, it won't "The company needs to design an application that captures AWS API calls". it only sends the "CreateImage API " event. We need to store the AWS API calls as well.

upvoted 1 times

cookieMr 1 year, 11 months ago

EventBridge (formerly CloudWatch Events) is a fully managed event bus service that allows you to monitor and respond to events within your AWS environment. By creating an EventBridge rule specifically for the CreateImage API call, you can easily detect and capture this event. Configuring the target as an SNS topic allows you to send an alert whenever a CreateImage API call occurs. This solution requires minimal operational overhead as EventBridge and SNS are fully managed services.

A. While using an Lambda to query CloudTrail logs and send an alert can achieve the desired outcome, it introduces additional operational overhead compared to using EventBridge and SNS directly.

B. Configuring CloudTrail with an SNS notification and using Athena to query on CreateImage API calls would require more setup and maintenance compared to using EventBridge and SNS.

D. Configuring an SQS FIFO queue as a target for CloudTrail logs and using a function to send an alert to an SNS topic adds unnecessary complexity to the solution and increases operational overhead. Using EventBridge and SNS directly is a simpler and more efficient approach.

upvoted 9 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #207

Topic 1

A company owns an asynchronous API that is used to ingest user requests and, based on the request type, dispatch requests to the appropriate microservice for processing. The company is using Amazon API Gateway to deploy the API front end, and an AWS Lambda function that invokes Amazon DynamoDB to store user requests before dispatching them to the processing microservices.

The company provisioned as much DynamoDB throughput as its budget allows, but the company is still experiencing availability issues and is losing user requests.

What should a solutions architect do to address this issue without impacting existing users?

- A. Add throttling on the API Gateway with server-side throttling limits.
- B. Use DynamoDB Accelerator (DAX) and Lambda to buffer writes to DynamoDB.
- C. Create a secondary index in DynamoDB for the table with the user requests.
- D. Use the Amazon Simple Queue Service (Amazon SQS) queue and Lambda to buffer writes to DynamoDB. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

pentium75 **Highly Voted** 11 months, 2 weeks ago

Selected Answer: D

A does not meet the "without impacting existing users" requirement
 B does not help with writing (DAX caches reads)
 C does not help with writing (index could increase read performance only)
 D decouples writing from front-end, which is acceptable because it is "an asynchronous API" anyway
 upvoted 14 times

nder **Highly Voted** 1 year, 9 months ago

Selected Answer: D

The key here is "Losing user requests" sqs messages will stay in the queue until it has been processed
 upvoted 12 times

lofzee Most Recent 6 months, 1 week ago

Selected Answer: D

D bro. Believe
upvoted 3 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: D

This solution can handle bursts of incoming requests more effectively and reduce the chances of losing requests due to DynamoDB capacity limitations. The Lambda can be configured to retrieve messages from the SQS and write them to DynamoDB at a controlled rate, allowing DynamoDB to handle the requests within its provisioned capacity. This approach provides resilience to spikes in traffic and ensures that requests are not lost during periods of high demand.

upvoted 5 times

cookieMr 1 year, 5 months ago

Selected Answer: D

This solution can handle bursts of incoming requests more effectively and reduce the chances of losing requests due to DynamoDB capacity limitations. The Lambda can be configured to retrieve messages from the SQS and write them to DynamoDB at a controlled rate, allowing DynamoDB to handle the requests within its provisioned capacity. This approach provides resilience to spikes in traffic and ensures that requests are not lost during periods of high demand.

A. It limits can help control the request rate, but it may lead to an increase in errors and affect the user experience. Throttling alone may not be sufficient to address the availability issues and prevent the loss of requests.

B. It can improve read performance but does not directly address the availability issues and loss of requests. It focuses on optimizing read operations rather than buffering writes.

C. It may help with querying the user requests efficiently, but it does not directly solve the availability issues or prevent the loss of requests. It is more focused on data retrieval rather than buffering writes.

upvoted 4 times

studynoplay 1 year, 6 months ago

Selected Answer: D

DAX is for reads
upvoted 4 times

smartegnine 1 year, 6 months ago

DAX is not ideal for the following types of applications:

Applications that require strongly consistent reads (or that cannot tolerate eventually consistent reads).

Applications that do not require microsecond response times for reads, or that do not need to offload repeated read activity from underlying tables.

Applications that are write-intensive, or that do not perform much read activity.

Applications that are already using a different caching solution with DynamoDB, and are using their own client-side logic for working with that caching solution.

upvoted 3 times

dark_firzen 1 year, 10 months ago

Selected Answer: D

D because SQS is the cheapest way. First 1,000,000 requests are free each month.

Question states: "The company provisioned as much DynamoDB throughput as its budget allows"

upvoted 4 times

Wajif 1 year, 11 months ago

Selected Answer: D

D is more likely to fix this problem as SQS queue has the ability to wait (buffer) for consumer to notify that the request or message has been processed.

upvoted 2 times

Buruguduystunstugudunstuy 1 year, 11 months ago

Selected Answer: D

To address the issue of lost user requests and improve the availability of the API, the solutions architect should use the Amazon Simple Queue Service (Amazon SQS) queue and Lambda to buffer writes to DynamoDB. Option D (correct answer)

By using an SQS queue and Lambda, the solutions architect can decouple the API front end from the processing microservices and improve the overall scalability and availability of the system. The SQS queue acts as a buffer, allowing the API front end to continue accepting user requests even if the processing microservices are experiencing high workloads or are temporarily unavailable. The Lambda function can then retrieve requests from the SQS queue and write them to DynamoDB, ensuring that all user requests are stored and processed. This approach allows the company to scale the processing microservices independently from the API front end, ensuring that the API remains available to users even during periods of high demand.

upvoted 7 times

alect096 1 year, 11 months ago

Selected Answer: B

I would go to B : <https://aws.amazon.com/es/blogs/database/amazon-dynamodb-accelerator-dax-a-read-throughwrite-through-cache-for-dynamodb/>

upvoted 1 times

ruqui 1 year, 5 months ago

That's wrong. The document you mentioned explained it very clearly:

"Whereas both read-through and write-through caches address read-heavy workloads, a write-back (or write-behind) cache is designed to address write-heavy workloads. Note that DAX is not a write-back cache currently"

upvoted 3 times

BENICE 1 year, 11 months ago

D is correct answer

upvoted 2 times

NikaCZ 1 year, 11 months ago

Selected Answer: D

D. Use the Amazon Simple Queue Service (Amazon SQS) queue and Lambda to buffer writes to DynamoDB.

upvoted 2 times

career360guru 1 year, 11 months ago

Selected Answer: D

Option D is right answer

upvoted 2 times

alexfk 1 year, 11 months ago

Why not B? DAX.

"When you're developing against DAX, instead of pointing your application at the DynamoDB endpoint, you point it at the DAX endpoint, and DAX handles the rest. As a read-through/write-through cache, DAX seamlessly intercepts the API calls that an application normally makes to DynamoDB so that both read and write activity are reflected in the DAX cache."

<https://aws.amazon.com/es/blogs/database/amazon-dynamodb-accelerator-dax-a-read-throughwrite-through-cache-for-dynamodb/>

upvoted 1 times

AgboolaKun 1 year, 7 months ago

It is not DAX because of the company's budget restriction associated with the DynamoDB. This is a requirement in the question. DynamoDB charges for DAX capacity by the hour and your DAX instances run with no long-term commitments. Please refer to: https://aws.amazon.com/dynamodb/pricing/provisioned/#.E2.80.A2_DynamoDB_Accelerator_.28DAX.29

upvoted 3 times

ruqui 1 year, 5 months ago

B is wrong because of this:

"Whereas both read-through and write-through caches address read-heavy workloads, a write-back (or write-behind) cache is designed to address write-heavy workloads. Note that DAX is not a write-back cache currently"

upvoted 2 times

akosigengen 2 years ago

yeah I though the answer is also DAX.

upvoted 1 times

leonnnn 2 years ago

Selected Answer: D

Using SQS should be the answer.

upvoted 4 times

nVizzz 2 years ago

Why not DAX? Could somebody explain?

upvoted 1 times

Buruguduystunstugudunstuy 1 year, 11 months ago

Using DynamoDB Accelerator (DAX) and Lambda to buffer writes to DynamoDB, may improve the write performance of the system, but it does not provide the same level of scalability and availability as using an SQS queue and Lambda.

Hence, Option B is incorrect.

upvoted 2 times

bmofo 2 years ago

key noted issue is "losing user requests" which is resolved with SQS

upvoted 6 times

Rameez1 2 years ago

DAX helps in reducing the read loads from DynamoDB, here we need a solution to handle write requests, which is well handled by SQS and Lamda to buffer writes on DynamoDB.

upvoted 5 times

jambajuice 2 years ago

Selected Answer: D

Answer d

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #208

Topic 1

A company needs to move data from an Amazon EC2 instance to an Amazon S3 bucket. The company must ensure that no API calls and no data are routed through public internet routes. Only the EC2 instance can have access to upload data to the S3 bucket.

Which solution will meet these requirements?

- A. Create an interface VPC endpoint for Amazon S3 in the subnet where the EC2 instance is located. Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access. **Most Voted**
- B. Create a gateway VPC endpoint for Amazon S3 in the Availability Zone where the EC2 instance is located. Attach appropriate security groups to the endpoint. Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.
- C. Run the nslookup tool from inside the EC2 instance to obtain the private IP address of the S3 bucket's service API endpoint. Create a route in the VPC route table to provide the EC2 instance with access to the S3 bucket. Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.
- D. Use the AWS provided, publicly available ip-ranges.json file to obtain the private IP address of the S3 bucket's service API endpoint. Create a route in the VPC route table to provide the EC2 instance with access to the S3 bucket. Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.

Correct Answer: A*Community vote distribution***Comments****SSASSWS** **Highly Voted** 2 years, 6 months ago**Selected Answer: A**

I think answer should be A and not B.
as we cannot "Attach a security groups to a gateway endpoint."
upvoted 41 times

A_New_Guy 2 years, 5 months ago

It's possible:

<https://aws.amazon.com/premiumsupport/knowledge-center/connect-s3-vpc-endpoint/>

upvoted 5 times

kruasan 2 years, 1 month ago

No, it's not

upvoted 3 times

smartegnine 2 years ago

Create a security group that allows the resources in your VPC to communicate with the endpoint network interfaces for the VPC endpoint. To ensure that tools such as the AWS CLI can make requests over HTTPS from resources in the VPC to the AWS service, the security group must allow inbound HTTPS traffic.

For Security groups, select the security groups to associate with the endpoint network interfaces for the VPC endpoint. By default, we associate the default security group for the VPC.

upvoted 2 times

slackbot 1 year, 9 months ago

this is valid for interface endpoint, not for gateway endpoint, which option B mentioned

upvoted 5 times

markw92 1 year, 11 months ago

Gateway endpoint must be used as a target in a route table does not use security groups.

upvoted 9 times

Iconique 1 year, 8 months ago

Go to console and test it yourself! With Interface Endpoint you can add security groups.

upvoted 2 times

elmyth 7 months, 1 week ago

interface VPC endpoint is A))))

upvoted 2 times

Buruguduystunstugudunstuy Highly Voted 2 years, 5 months ago

Selected Answer: B

The correct solution to meet the requirements is Option B. A gateway VPC endpoint for Amazon S3 should be created in the Availability Zone where the EC2 instance is located. This will allow the EC2 instance to access the S3 bucket directly, without routing through the public internet. The endpoint should also be configured with appropriate security groups to allow access to the S3 bucket. Additionally, a resource policy should be attached to the S3 bucket to only allow the EC2 instance's IAM role for access.

upvoted 36 times

JA2018 6 months, 3 weeks ago

check this out: <https://www.examtopycs.com/discussions/amazon/view/83857-exam-aws-certified-solutions-architect-associate-saa-c02/>

Selected answer: B

So which is the correct answer?!!!!

upvoted 1 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option A is incorrect because an interface VPC endpoint for Amazon S3 would not provide a direct connection between the EC2 instance and the S3 bucket.

Option C is incorrect because using the nslookup tool to obtain the private IP address of the S3 bucket's service API endpoint would not provide a secure connection between the EC2 instance and the S3 bucket.

Option D is incorrect because using the ip-ranges.json file to obtain the private IP address of the S3 bucket's service API

endpoint is not a secure method to connect the EC2 instance to the S3 bucket.
upvoted 4 times

ChrisG1454 2 years, 3 months ago

There are two types VPC Endpoint:

Gateway endpoint
Interface endpoint

A Gateway endpoint:

- 1) Helps you to securely connect to Amazon S3 and DynamoDB
- 2) Endpoint serves as a target in your route table for traffic
- 3) Provide access to endpoint (endpoint, identity and resource policies)

An Interface endpoint:

- 1) Help you to securely connect to AWS services EXCEPT FOR Amazon S3 and DynamoDB
- 2) Powered by PrivateLink (keeps network traffic within AWS network)
- 3) Needs a elastic network interface (ENI) (entry point for traffic)

upvoted 35 times

slackbot 1 year, 9 months ago

interface endpoint exists for S3 as well
upvoted 10 times

mhmt4438 2 years, 5 months ago

An interface VPC endpoint does provide a direct connection between the EC2 instance and the S3 bucket. It enables private communication between instances in your VPC and resources in other services without requiring an internet gateway, a NAT device, or a VPN connection.

Option A , which recommends creating an interface VPC endpoint for Amazon S3 in the subnet where the EC2 instance is located and attaching a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access, is the correct solution for the given scenario. It meets the requirement to ensure that no API calls and no data are routed through public internet routes and that only the EC2 instance can have access to upload data to the S3 bucket.

upvoted 7 times

Omok 2 years, 4 months ago

In support, see <https://docs.aws.amazon.com/AmazonS3/latest/userguide/privatelink-interface-endpoints.html#types-of-vpc-endpoints-for-s3>
upvoted 8 times

TEC65 Most Recent 1 month, 2 weeks ago

Selected Answer: B

Here's why this option is suitable:

Gateway VPC Endpoint for Amazon S3: A gateway VPC endpoint allows you to privately connect your VPC to supported AWS services, such as Amazon S3, without requiring an internet gateway, NAT device, VPN connection, or AWS Direct Connect connection. Traffic between your VPC and the service does not leave the Amazon network.

Resource Policy: By attaching a resource policy to the S3 bucket that only allows access from the EC2 instance's IAM role, you ensure that only authorized instances can access the bucket.

Here's why the other options are not as suitable:

A. Create an interface VPC endpoint for Amazon S3:

Amazon S3 uses gateway VPC endpoints, not interface VPC endpoints. Interface VPC endpoints are used for services that are not supported by gateway endpoints.

upvoted 1 times

7dcef09 1 month, 4 weeks ago

Selected Answer: B

Interface Endpoint = AWS PrivateLink

Gateway Endpoint = Supports S3 and DynamoDB

upvoted 1 times

upvoted 1 times

nadeerm 2 months, 4 weeks ago

Reasons for B:

Simplicity: Gateway VPC endpoints are easier to configure and manage compared to interface VPC endpoints.

Cost: Gateway VPC endpoints are free to use, whereas interface VPC endpoints incur costs for each hour and per GB of data processed.

Performance: Gateway VPC endpoints are highly scalable and optimized for S3 and DynamoDB traffic.

upvoted 1 times

chir827389_3 3 months, 4 weeks ago

Selected Answer: A

gateway endpoint does not have any SG.

upvoted 1 times

certifiedlegend 3 months, 4 weeks ago

Selected Answer: B

A VPC gateway endpoint is primarily used for accessing specific AWS services like Amazon S3 and DynamoDB privately within your VPC by specifying a route in your route table, while a VPC interface endpoint offers more flexible connectivity to a wider range of AWS services through AWS PrivateLink, allowing access from both within your VPC and from other VPCs using peering or Transit Gateways, typically with a dedicated private IP address within your VPC network; essentially, gateway endpoints are simpler and often free, while interface endpoints provide greater control and may incur additional costs depending on usage. Gateway endpoints are ideal for simple private access to S3 and DynamoDB, while interface endpoints are better suited for more complex scenarios where you need to access various AWS services from different VPCs or on-premises networks.

upvoted 1 times

DeliKadir 4 months ago

Selected Answer: A

Even though Amazon S3 typically uses a Gateway VPC Endpoint, AWS now provides Interface VPC Endpoints for Amazon S3 as well via AWS PrivateLink....and "Attach a security groups to a gateway endpoint." makes the option B - false

upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

Selected Answer: A

B will be correct if removing the "Attach appropriate security group to the endpoint" and gateway point is free of charge.

upvoted 1 times

MaxMingxing 5 months, 2 weeks ago

Selected Answer: B

Option A uses an interface VPC endpoint, which is typically used for services that require a private IP address within your VPC. For S3, a gateway VPC endpoint is more appropriate and cost-effective¹

upvoted 1 times

hilker1983 5 months, 2 weeks ago

Selected Answer: B

Interface VPC Endpoint: Interface endpoints are generally used for other AWS services and do not provide the same direct access optimization for S3 that gateway endpoints do.

upvoted 1 times

LizavetaD 5 months, 2 weeks ago

Selected Answer: B

In this context, a Gateway VPC Endpoint is the correct choice for S3, as it provides direct, private access to S3 and routes traffic internally within the AWS network

upvoted 1 times

dipenich 5 months, 3 weeks ago

Selected Answer: B

Why Option B is Correct:

Gateway VPC Endpoint for Amazon S3:

Gateway VPC endpoint for Amazon S3:

S3 does not require an interface endpoint; it uses a gateway VPC endpoint.

Gateway endpoints ensure that requests stay within the AWS network, meeting the requirement of no public internet routes.

Resource Policy:

Attaching a bucket policy ensures that only the EC2 instance's IAM role has access to the bucket.

Appropriate Security Controls:

The gateway endpoint can be secured further using policies and security group configurations.

upvoted 1 times

ARV14 5 months, 4 weeks ago

Selected Answer: A

It is A. Gateway endpoint route table. Interface endpoint uses security groups and private link. See aws video

<https://youtu.be/TqApkvJx5hw?si=9Gpk3V7OcPU6MVJI>

upvoted 1 times

SteveNguyen 6 months ago

Selected Answer: B

gateway VPC endpoint is used for connect from VPC to S3 and DynamoDB

upvoted 1 times

FlyingHawk 6 months, 2 weeks ago

Selected Answer: A

You might initially lean toward Option B since a gateway endpoint is generally the preferred solution for EC2 instances in the same region. It's cost-effective (free), performance-optimized, and simpler to configure. However, for the exam, technical precision in details is crucial.

Option B is incorrect because:

Gateway VPC endpoints are created at the VPC level, not at the Availability Zone level.

Security groups must be attached to the EC2 instance (the source service) to allow outbound traffic using the prefix list associated with the gateway endpoint.

Due to these reasons, Option A is the correct answer. While not ideal in real-world scenarios for EC2 instances in the same region, it is technically accurate and satisfies the exam's requirements.

upvoted 4 times

Garryg 6 months, 2 weeks ago

Selected Answer: B

correct ans is B

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #209

Topic 1

A solutions architect is designing the architecture of a new application being deployed to the AWS Cloud. The application will run on Amazon EC2 On-Demand Instances and will automatically scale across multiple Availability Zones. The EC2 instances will scale up and down frequently throughout the day. An Application Load Balancer (ALB) will handle the load distribution. The architecture needs to support distributed session data management. The company is willing to make changes to code if needed.

What should the solutions architect do to ensure that the architecture supports distributed session data management?

- A. Use Amazon ElastiCache to manage and store session data. **Most Voted**
- B. Use session affinity (sticky sessions) of the ALB to manage session data.
- C. Use Session Manager from AWS Systems Manager to manage the session.
- D. Use the GetSessionToken API operation in AWS Security Token Service (AWS STS) to manage the session.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Buruguduystunstugudunstuy **Highly Voted** 1 year, 11 months ago

Selected Answer: A

The correct answer is A. Use Amazon ElastiCache to manage and store session data.

In order to support distributed session data management in this scenario, it is necessary to use a distributed data store such as Amazon ElastiCache. This will allow the session data to be stored and accessed by multiple EC2 instances across multiple Availability Zones, which is necessary for a scalable and highly available architecture.

Option B, using session affinity (sticky sessions) of the ALB, would not be sufficient because this would only allow the session data to be stored on a single EC2 instance, which would not be able to scale across multiple Availability Zones.

Options C and D, using Session Manager and the GetSessionToken API operation in AWS STS, are not related to session data management and would not be appropriate solutions for this scenario.

upvoted 34 times

TariqKipkemei **Highly Voted** 1 year, 2 months ago

Selected Answer: A

Yap agree with go you guys, this is one of the use cases for Amazon ElastiCache.

It was designed to store ephemeral session data to quickly personalize gaming, e-commerce, social media, and online applications with microsecond response times.

<https://aws.amazon.com/elasticache/#:~:text=Store-,ephemeral,-session%20data%20to>

upvoted 5 times

Cloud_A 11 months ago

<https://aws.amazon.com/elasticache/#:~:text=Store%20ephemeral%20session%20data%20to%20quickly%20personalize%20gaming%2C%20e%2Dcommerce%2C%20social%20media%2C%20and%20online%20applications%20with%20microsecond%20response%20times>.

upvoted 2 times

Hkayne **Most Recent** 7 months, 1 week ago

Selected Answer: A

A is the correct answer because it allows to manage distributed sessions

upvoted 3 times

awsgeek75 11 months, 1 week ago

Selected Answer: A

A is in scope of question as company is willing to make code changes.

B would have been correct if no code changes were allowed and scaling could be compromised.

C is wrong technology (cloud management)

D is also wrong technology (AWS IAM or account management).

upvoted 5 times

Michael_Li 1 year ago

A is correct

B is not correct as session affinity allow web user stick to a EC2 instance for a period time, that EC2 could go down then the session data will lost, so doesn't fit this use case

C is wrong as Session Manager is for admins users to manage EC2 CLI access, it's not for web end users

D is wrong as GetSessionToken API is for use case such as you need to grant user access to a S3 bucket with customized code

upvoted 3 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: A

The correct answer is A. Use Amazon ElastiCache to manage and store session data.

upvoted 2 times

cookieMr 1 year, 5 months ago

Selected Answer: A

ElastiCache is a managed in-memory data store service that is well-suited for managing session data in a distributed architecture. It provides high-performance, scalable, and durable storage for session data, allowing multiple EC2 instances to access and share session data seamlessly. By using ElastiCache, the application can offload the session management workload from the EC2 instances and leverage the distributed caching capabilities of ElastiCache for improved scalability and performance.

Option B, using session affinity (sticky sessions) of the ALB, is not the best choice for distributed session data management because it ties each session to a specific EC2 instance. As the instances scale up and down frequently, it can lead to uneven load distribution and may not provide optimal scalability.

Options C and D are not applicable for managing session data. AWS Systems Manager's Session Manager is primarily used for secure remote shell access to EC2 instances, and the AWS STS GetSessionToken API operation is used for temporary security credentials and not session data management.

upvoted 4 times

cookieMr 1 year, 5 months ago

ElastiCache is a managed in-memory data store service that is well-suited for managing session data in a distributed architecture. It provides high-performance, scalable, and durable storage for session data, allowing multiple EC2 instances to access and share session data seamlessly. By using ElastiCache, the application can offload the session management workload from the EC2 instances and leverage the distributed caching capabilities of ElastiCache for improved scalability and performance.

Option B, using session affinity (sticky sessions) of the ALB, is not the best choice for distributed session data management because it ties each session to a specific EC2 instance. As the instances scale up and down frequently, it can lead to uneven load distribution and may not provide optimal scalability.

Options C and D are not applicable for managing session data. AWS Systems Manager's Session Manager is primarily used for secure remote shell access to EC2 instances, and the AWS STS GetSessionToken API operation is used for temporary security credentials and not session data management.

upvoted 4 times

Abrar2022 1 year, 5 months ago

Selected Answer: A

A. Use Amazon ElastiCache to manage and store session data.

- Correct. - Session data is managed at the application-layer, and a distributed cache should be used

B. Use session affinity (sticky sessions) of the ALB to manage session data.

- Wrong. This tightly couples the individual EC2 instances to the session data, and requires additional logic in the ALB. When scale-in happens, the session data stored on individual EC2 instances is destroyed

upvoted 2 times

techhb 1 year, 10 months ago

Selected Answer: A

correct answer is A as instance are getting up and down.

upvoted 2 times

inseong 1 year, 11 months ago

□ □ 210□□ □□ ..?

upvoted 1 times

noche 1 year, 9 months ago

<https://www.examtopycs.com/discussions/amazon/view/94992-exam-aws-certified-solutions-architect-associate-saa-c03/>

□ □ □ □

upvoted 1 times

NikaCZ 1 year, 11 months ago

Selected Answer: A

Amazon ElastiCache to manage and store session data.

upvoted 2 times

k1kavi1 1 year, 11 months ago

Selected Answer: A

<https://www.examtopycs.com/discussions/amazon/view/46412-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Shasha1 1 year, 12 months ago

A
Amazon ElastiCache to manage and store session data. This solution will allow the application to automatically scale across multiple Availability Zones without losing session data, as the session data will be stored in a cache that is accessible from any EC2 instance. Additionally, using Amazon ElastiCache will enable the company to easily manage and scale the cache as needed, without requiring any changes to the application code. Option C is not correct because, Session Manager from AWS Systems Manager will not provide the necessary support for distributed session data management. Session Manager is a tool for managing and tracking sessions on EC2 instances, but it does not provide a mechanism for storing and managing session data in a distributed environment.

upvoted 4 times

Tela0 2 years ago

better justification found here...

<https://www.examtopycs.com/discussions/amazon/view/46412-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 4 times

kmaneith 2 years ago

why not C?

upvoted 1 times

leonnnn 2 years ago

Selected Answer: A

ALB sticky session can keep request accessing to the same backend application. But it says "distributed session management" and company "will to change code", so I think A is better

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #210

Topic 1

A company offers a food delivery service that is growing rapidly. Because of the growth, the company's order processing system is experiencing scaling problems during peak traffic hours. The current architecture includes the following:

- A group of Amazon EC2 instances that run in an Amazon EC2 Auto Scaling group to collect orders from the application
- Another group of EC2 instances that run in an Amazon EC2 Auto Scaling group to fulfill orders

The order collection process occurs quickly, but the order fulfillment process can take longer. Data must not be lost because of a scaling event.


A solutions architect must ensure that the order collection process and the order fulfillment process can both scale properly during peak traffic hours. The solution must optimize utilization of the company's AWS resources.

Which solution meets these requirements?

- A. Use Amazon CloudWatch metrics to monitor the CPU of each instance in the Auto Scaling groups. Configure each Auto Scaling group's minimum capacity according to peak workload values.
- B. Use Amazon CloudWatch metrics to monitor the CPU of each instance in the Auto Scaling groups. Configure a CloudWatch alarm to invoke an Amazon Simple Notification Service (Amazon SNS) topic that creates additional Auto Scaling groups on demand.
- C. Provision two Amazon Simple Queue Service (Amazon SQS) queues: one for order collection and another for order fulfillment. Configure the EC2 instances to poll their respective queue. Scale the Auto Scaling groups based on notifications that the queues send.
- D. Provision two Amazon Simple Queue Service (Amazon SQS) queues: one for order collection and another for order fulfillment. Configure the EC2 instances to poll their respective queue. Create a metric based on a backlog per instance calculation. Scale the Auto Scaling groups based on this metric. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

TungPham Highly Voted 2 years, 3 months ago

Selected Answer: D

When the backlog per instance reaches the target value, a scale-out event will happen. Because the backlog per instance is already 150 messages (1500 messages / 10 instances), your group scales out, and it scales out by five instances to maintain proportion to the target value.

Backlog per instance: To calculate your backlog per instance, start with the `ApproximateNumberOfMessages` queue attribute to determine the length of the SQS queue (number of messages available for retrieval from the queue). Divide that number by the fleet's running capacity, which for an Auto Scaling group is the number of instances in the `InService` state, to get the backlog per instance.

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-using-sqs-queue.html>

upvoted 16 times

n43u435b543ht2b Highly Voted 1 year, 10 months ago

Selected Answer: D

C is incorrect as scaling based on the number of "notifications" doesn't make logical sense. This means that both the order collection and fulfilment instances would scale in parallel, but they have clearly said that the collection is processing quickly while the fulfilment is struggling. Therefore, we should scale the pool when there is a backlog building in a respective queue - not just based on the number of incoming requests.

upvoted 13 times

jerry1 Most Recent 2 months, 2 weeks ago

Selected Answer: C

I can understand you guy's explanation on D (the burden is on one queue only not both) but i dont understand how D imply that the metric is monitoring on on side

upvoted 1 times

satyaamm 3 months, 3 weeks ago

Selected Answer: D

Since only the order fulfillment requires scaling capacity so D is the most suitable here.

upvoted 1 times

jatric 11 months ago

Selected Answer: C

both have their own queue. Instance processing order will be scale up based on the queue length that collect messages that collected by other queue.

upvoted 1 times

lofzee 1 year ago

Selected Answer: D

not C as the questions state that only one system is struggling, so C doesnt really solve the problem. D does.

upvoted 2 times

Uzbekistan 1 year, 2 months ago

Selected Answer: D

Decoupling with Amazon SQS: By using Amazon SQS queues for order collection and order fulfillment, the system can decouple the components, ensuring that orders are not lost, even during scaling events. Orders are queued up and processed in a reliable and scalable manner.

Scalability Based on Queue Backlog: By creating a metric based on the backlog per instance calculation, the system can monitor the workload of each instance in the Auto Scaling groups. This allows for dynamic scaling based on the workload, ensuring that additional instances are launched when the backlog increases and terminated when the backlog decreases. Optimization of AWS Resources: This solution optimizes the utilization of AWS resources by dynamically scaling the Auto Scaling groups based on the actual workload, preventing over-provisioning or under-provisioning of instances. It ensures that the system can handle peak traffic efficiently without incurring unnecessary costs.

upvoted 2 times

bujuman 1 year, 4 months ago

Selected Answer: D

D is the most appropriate response base on <https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-using-sqs-queue.html>

upvoted 2 times

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

D. Provision two Amazon Simple Queue Service (Amazon SQS) queues: one for order collection and another for order fulfillment. Configure the EC2 instances to poll their respective queue. Create a metric based on a backlog per instance calculation. Scale the Auto Scaling groups based on this metric.

upvoted 5 times

argl1995 1 year, 11 months ago

SQS auto-scales by default so I don't think we need to mention it explicitly. Option D should be correct.

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: D

A. This approach focuses solely on CPU utilization, which may not accurately reflect the scaling needs of the order collection and fulfillment processes. It does not address the need for decoupling and reliable message processing.

B. While this approach incorporates alarms to trigger additional Auto Scaling groups, it lacks the decoupling and reliable message processing provided by using SQS queues. It may lead to inefficient scaling and potential data loss.

C. Although using SQS queues is a step in the right direction, scaling solely based on queue notifications may not provide optimal resource utilization. It does not consider the backlog per instance and does not allow for fine-grained control over scaling.

Overall, option D, which involves using SQS queues for order collection and fulfillment, creating a metric based on backlog per instance calculation, and scaling the Auto Scaling groups accordingly, is the most suitable solution to address the scaling problems while optimizing resource utilization and ensuring reliable message processing.

upvoted 5 times

studynoplay 2 years ago

Selected Answer: D

C is incorrect. "based on notifications that the queues send" SQS does not send notification

upvoted 4 times

mandragon 2 years, 1 month ago

Selected Answer: C

D is not correct because it requires more operational overhead and complexity than option C which is simpler and more cost-effective. It uses the existing queue metrics that are provided by Amazon SQS and does not require creating or publishing any custom metrics. You can use target tracking scaling policies to automatically maintain a desired backlog per instance ratio without having to calculate or monitor it yourself.

upvoted 2 times

pentium75 1 year, 5 months ago

"You can use target tracking scaling policies" but you don't with option C. What is "scaling based on notifications that the queues send"? Where do they send these notifications to?

upvoted 1 times

JayBee65 2 years, 4 months ago

Selected Answer: D

Scale based on queue length

upvoted 3 times

Rudraman 2 years, 4 months ago

answer is D.

read question again

upvoted 3 times

LuckyAro 2 years, 4 months ago

Selected Answer: D

The number of instances in your Auto Scaling group can be driven by how long it takes to process a message and the acceptable amount of latency (queue delay).

The solution is to use a backlog per instance metric with the target value being the acceptable backlog per instance to maintain.

upvoted 2 times

Aseem8888 2 years, 4 months ago

Selected Answer: D

D is correct

upvoted 2 times