

Option A - FSx for NetApp ONTAP with ALL tiering policy would not provide fast enough storage tier for sub-millisecond latency. HDD tiers have higher latency.

Option C - FSx for Lustre with HDD storage would not provide the throughput, IOPS or low latency needed.

Option D - FSx for NetApp ONTAP with NONE tiering policy would require much more expensive SSD storage to meet requirements, increasing cost.

upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: B

I vote B

upvoted 2 times

AlmeroSenior 1 year, 9 months ago

Selected Answer: B

FSX Lustre is 1000mbs per TB provisioned and we have 8TBs so gives us 8GBs . The netapp FSX appears a hard limit of 4gbs .

<https://aws.amazon.com/fsx/lustre/faqs/?nc=sn&loc=5>

<https://aws.amazon.com/fsx/netapp-ontap/faqs/>

upvoted 6 times

LuckyAro 1 year, 9 months ago

Selected Answer: B

B is the best choice as it utilizes Amazon S3 for data storage, which is cost-effective and durable, and Amazon FSx for Lustre for high-performance file storage, which provides the required sub-millisecond latencies and minimum throughput of 6 GBps. Additionally, the option to import and export data to and from Amazon S3 makes it easier to manage and move data between the two services.

B is the best option as it meets the performance requirements for sub-millisecond latencies and a minimum throughput of 6 GBps.

upvoted 2 times

everfly 1 year, 9 months ago

Selected Answer: B

Amazon FSx for Lustre provides fully managed shared storage with the scalability and performance of the popular Lustre file system. It can deliver sub-millisecond latencies and hundreds of gigabytes per second of throughput.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #300

Topic 1

A company needs to migrate a legacy application from an on-premises data center to the AWS Cloud because of hardware capacity constraints. The application runs 24 hours a day, 7 days a week. The application's database storage continues to grow over time.

What should a solutions architect do to meet these requirements MOST cost-effectively?

- A. Migrate the application layer to Amazon EC2 Spot Instances. Migrate the data storage layer to Amazon S3.
- B. Migrate the application layer to Amazon EC2 Reserved Instances. Migrate the data storage layer to Amazon RDS On-Demand Instances.
- C. Migrate the application layer to Amazon EC2 Reserved Instances. Migrate the data storage layer to Amazon Aurora Reserved Instances. **Most Voted**
- D. Migrate the application layer to Amazon EC2 On-Demand Instances. Migrate the data storage layer to Amazon RDS Reserved Instances.

Correct Answer: C*Community vote distribution***Comments****LuckyAro** **Highly Voted** 1 year, 9 months ago**Selected Answer: C**

Amazon EC2 Reserved Instances allow for significant cost savings compared to On-Demand instances for long-running, steady-state workloads like this one. Reserved Instances provide a capacity reservation, so the instances are guaranteed to be available for the duration of the reservation period.

Amazon Aurora is a highly scalable, cloud-native relational database service that is designed to be compatible with MySQL and PostgreSQL. It can automatically scale up to meet growing storage requirements, so it can accommodate the application's database storage needs over time. By using Reserved Instances for Aurora, the cost savings will be significant over the long term.

upvoted 22 times

NolaHOLA **Highly Voted** 1 year, 9 months ago

Option B based on the fact that the DB storage will continue to grow, so on-demand will be a more suitable solution

upvoted 15 times

NolaHOla 1 year, 9 months ago

Since the application's database storage is continuously growing over time, it may be difficult to estimate the appropriate size of the Aurora cluster in advance, which is required when reserving Aurora.

In this case, it may be more cost-effective to use Amazon RDS On-Demand Instances for the data storage layer. With RDS On-Demand Instances, you pay only for the capacity you use and you can easily scale up or down the storage as needed.

upvoted 5 times

hristni0 1 year, 6 months ago

Answer is C. From Aurora Reserved Instances documentation:

If you have a DB instance, and you need to scale it to larger capacity, your reserved DB instance is automatically applied to your scaled DB instance. That is, your reserved DB instances are automatically applied across all DB instance class sizes. Size-flexible reserved DB instances are available for DB instances with the same AWS Region and database engine.

upvoted 2 times

Joxtat 1 year, 9 months ago

The Answer is C.

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.AuroraMySQL.html>

upvoted 2 times

pentium75 11 months, 2 weeks ago

Database STORAGE will grow, not performance need (and required instance size).

upvoted 4 times

zdi561 Most Recent 4 months ago

Selected Answer: B

RLs are best for steady-state database workloads.

upvoted 1 times

hro 8 months, 2 weeks ago

cost-effectively - the answer is C.

The application runs 24 hours a day, 7 days a week. The application's database storage continues to grow over time.

upvoted 3 times

MrPCarrot 9 months, 2 weeks ago

Answer is C: Amazon EC2 Reserved Instances and Amazon Aurora Reserved Instances = less expensive than RDS.

upvoted 3 times

andyngh86 10 months, 2 weeks ago

Amazon Aurora reserved instances is used for the work load on predictable, so answer should be B

upvoted 1 times

Priyapani 10 months, 4 weeks ago

Selected Answer: B

I think it's B as database storage will grow

upvoted 2 times

awsgeek75 10 months, 3 weeks ago

Application runs 24x7 which means database is also used 24x7. The storage will grow and RDS On-Demand does not have auto-grow storage. You have to configure a storage size for RDS which means it will eventually run out of space. RDS On-Demand just scales CPU, not storage.

Aurora has no storage limitation and can scale storage according to need which is what is required here

upvoted 4 times

Mikado211 1 year ago

Selected Answer: C

24/7 forbids spot instances , so A is excluded
Cost efficiency require reserved instances , so D is excluded
Between RDS and Aurora, Aurora is less expensive thanks to the reserved instance, so B is finally excluded

Answer is C
upvoted 4 times

cciesam 1 year, 1 month ago

Selected Answer: B

I hope it should be B considering Database growth
upvoted 1 times

pentium75 11 months, 2 weeks ago

Reserved instance applies to the DB instance size (CPU, RAM etc.), not storage.
upvoted 3 times

Wayne23Fang 1 year, 2 months ago

My research concludes that From pure price point of view Aurora Reserved might/ usually be slightly more expensive than On-demand RDS. But RDS has less Operation overhead. For the 24x7 nature, I would vote C. But for pure cost-effective, B is less costly.
upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

This option involves migrating the application layer to Amazon EC2 Reserved Instances and migrating the data storage layer to Amazon Aurora Reserved Instances. Amazon EC2 Reserved Instances provide a significant discount (up to 75%) compared to On-Demand Instance pricing, making them a cost-effective choice for applications that have steady state or predictable usage. Similarly, Amazon Aurora Reserved Instances provide a significant discount (up to 69%) compared to On-Demand Instance pricing.
upvoted 2 times

ajchi1980 1 year, 5 months ago

Selected Answer: C

To meet the requirements of migrating a legacy application from an on-premises data center to the AWS Cloud in a cost-effective manner, the most suitable option would be:

C. Migrate the application layer to Amazon EC2 Reserved Instances. Migrate the data storage layer to Amazon Aurora Reserved Instances.

Explanation:

Migrating the application layer to Amazon EC2 Reserved Instances allows you to reserve EC2 capacity in advance, providing cost savings compared to On-Demand Instances. This is especially beneficial if the application runs 24/7.

Migrating the data storage layer to Amazon Aurora Reserved Instances provides cost optimization for the growing database storage needs. Amazon Aurora is a fully managed relational database service that offers high performance, scalability, and cost efficiency.

upvoted 2 times

cpn 1 year, 6 months ago

nnascncscnknkckl
upvoted 1 times

TariqKipkemei 1 year, 7 months ago

Answer is C
upvoted 2 times

QuangPham810 1 year, 7 months ago

Answer is C. Refer
https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/USER_WorkingWithReservedDBInstances.html => Size-flexible reserved DB instances

upvoted 2 times

Abhineet9148232 1 year, 9 months ago

Selected Answer: C

C: With Aurora Serverless v2, each writer and reader has its own current capacity value, measured in ACUs. Aurora Serverless v2 scales a writer or reader up to a higher capacity when its current capacity is too low to handle the load. It scales the writer or reader down to a lower capacity when its current capacity is higher than needed.

This is sufficient to accommodate the growing data changes.

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-serverless-v2.how-it-works.html#aurora-serverless-v2.how-it-works.scaling>

upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: C

Typically Amazon RDS cost less than Aurora. But here, it's Aurora reserved.

upvoted 2 times

djgodzilla 11 months, 1 week ago

although agree and AWS wants you to choose Answer C. You can't convince a cloud accounting analyst that Aurora is cheaper than RDS. no matter what

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #301

Topic 1

A university research laboratory needs to migrate 30 TB of data from an on-premises Windows file server to Amazon FSx for Windows File Server. The laboratory has a 1 Gbps network link that many other departments in the university share.


The laboratory wants to implement a data migration service that will maximize the performance of the data transfer. However, the laboratory needs to be able to control the amount of bandwidth that the service uses to minimize the impact on other departments. The data migration must take place within the next 5 days.

Which AWS solution will meet these requirements?

- A. AWS Snowcone
- B. Amazon FSx File Gateway
- C. AWS DataSync **Most Voted**
- D. AWS Transfer Family

Correct Answer: C

Community vote distribution



Comments

kruasan **Highly Voted** 2 years, 1 month ago

Selected Answer: C

AWS DataSync is a data transfer service that can copy large amounts of data between on-premises storage and Amazon FSx for Windows File Server at high speeds. It allows you to control the amount of bandwidth used during data transfer.

- DataSync uses agents at the source and destination to automatically copy files and file metadata over the network. This optimizes the data transfer and minimizes the impact on your network bandwidth.
- DataSync allows you to schedule data transfers and configure transfer rates to suit your needs. You can transfer 30 TB within 5 days while controlling bandwidth usage.
- DataSync can resume interrupted transfers and validate data to ensure integrity. It provides detailed monitoring and reporting on the progress and performance of data transfers.

upvoted 27 times

kruasan 2 years, 1 month ago

Option A - AWS Snowcone is more suitable for physically transporting data when network bandwidth is limited. It would not

Option A - AWS Snowcone is more suitable for physically transporting data when network bandwidth is limited. It would not complete the transfer within 5 days.
Option B - Amazon FSx File Gateway only provides access to files stored in Amazon FSx and does not perform the actual data migration from on-premises to FSx.
Option D - AWS Transfer Family is for transferring files over FTP, FTPS and SFTP. It may require scripting to transfer 30 TB and monitor progress, and lacks bandwidth controls.

upvoted 16 times

Michal_L_95 Highly Voted 2 years, 2 months ago

Selected Answer: C

As read a little bit, I assume that B (FSx File Gateway) requires a little bit more configuration rather than C (DataSync). From Stephane Maarek course explanation about DataSync:
An online data transfer service that simplifies, automates, and accelerates copying large amounts of data between on-premises storage systems and AWS Storage services, as well as between AWS Storage services.

You can use AWS DataSync to migrate data located on-premises, at the edge, or in other clouds to Amazon S3, Amazon EFS, Amazon FSx for Windows File Server, Amazon FSx for Lustre, Amazon FSx for OpenZFS, and Amazon FSx for NetApp ONTAP.

upvoted 12 times

MatAlves Most Recent 9 months ago

Selected Answer: C

Even if we allocate 60% of total bandwidth for the transfer, that would take 5d2h. Considering that "many other departments in the university share", that wouldn't be feasible.

Ref. <https://expedient.com/knowledgebase/tools-and-calculators/file-transfer-time-calculator/>

On the other hand, snowcone isn't also a great option, because "you will receive the Snowcone device in approximately 4-6 days".

Ref.

<https://aws.amazon.com/snowcone/faqs/#:~:text=You%20will%20receive%20the%20Snowcone,console%20for%20each%20Snowcone%20device.>

upvoted 2 times

MatAlves 9 months ago

Very unreasonable scenario. The least "bad" option is C, but that will definitely affect users in real production environments.

upvoted 2 times

Cyberkayu 1 year, 5 months ago

Selected Answer: A

Snow cone can support up to 8TB for HDD and 15TB for each SSD devices. Shipped within 4-6 days. Data migration can begin on next 5 days.

Does not use any amount of bandwidth and impact the production network. Device came with 1G and 10G Base-T Ethernet port. That's the Maximum performance in data transfer. defined in the question.

upvoted 2 times

JA2018 6 months, 3 weeks ago

Key in STEM:

The data migration must take place within the next 5 days.

upvoted 1 times

JA2018 6 months, 3 weeks ago

Snowcone it's just 15TB per SSD and it takes 4-6 business days to deliver to the customer's location....

It could takes another 4-6 business days to return the snowcone device back to AWS (to transfer data to the designated S3 bucket) after you completed the on-prems data transfer to the device.

upvoted 1 times

AZ_Master 1 year, 6 months ago

Selected Answer: C

Bandwidth control = Data Sync

<https://docs.aws.amazon.com/datasync/latest/userguide/configure-bandwidth.html>

upvoted 5 times

Ruffyt 1 year, 6 months ago

Bandwidth Optimization and Control

Transferring hot or cold data should not impede your business. DataSync is equipped with granular controls to optimize bandwidth consumptions. Throttle transfer speeds up to 10 Gbps during off hours and set limits when network availability is needed elsewhere

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

C. AWS DataSync

upvoted 2 times

Nikki013 1 year, 9 months ago

Selected Answer: C

<https://aws.amazon.com/datasync/features/>

upvoted 2 times

Yousuf_Ibrahim 1 year, 8 months ago

Bandwidth Optimization and Control

Transferring hot or cold data should not impede your business. DataSync is equipped with granular controls to optimize bandwidth consumptions. Throttle transfer speeds up to 10 Gbps during off hours and set limits when network availability is needed elsewhere.

upvoted 2 times

jayce5 1 year, 12 months ago

Selected Answer: C

"Amazon FSx File Gateway" is for storing data, not for migrating. So the answer should be C.

upvoted 3 times

ACloud_Guru15 1 year, 7 months ago

Thanks for the explanation

upvoted 2 times

shanwford 2 years, 2 months ago

Selected Answer: C

Snowcone too small and delivery time too long. With DataSync you can set bandwidth limits - so this is a fine solution.

upvoted 3 times

MaxMa 2 years, 2 months ago

Why not B?

upvoted 1 times

Guru4Cloud 1 year, 8 months ago

Transferring will be much longer term rather than 5 days as required.

upvoted 2 times

AlessandraSAA 2 years, 3 months ago

A not possible because Snowcone it's just 8TB and it takes 4-6 business days to deliver

B why cannot be <https://aws.amazon.com/storagegateway/file/fsx/>?

C I don't really get this

D cannot be because not compatible - <https://aws.amazon.com/aws-transfer-family/>

upvoted 1 times

pentium75 1 year, 5 months ago

With B you cannot "control the amount of bandwidth that the service uses", while C does exactly what is required here.

upvoted 2 times

Steve_4542636 2 years, 3 months ago

Selected Answer: C

Voting C

upvoted 2 times

Bhawesh 2 years, 3 months ago

Selected Answer: C

C. - DataSync is Correct.

A. Snowcone is incorrect. The question says data migration must take place within the next 5 days. AWS says: If you order, you will receive the Snowcone device in approximately 4-6 days.

upvoted 3 times

LuckyAro 2 years, 3 months ago

Selected Answer: C

DataSync can be used to migrate data between on-premises Windows file servers and Amazon FSx for Windows File Server with its compatibility for Windows file systems.

The laboratory needs to migrate a large amount of data (30 TB) within a relatively short timeframe (5 days) and limit the impact on other departments' network traffic. Therefore, AWS DataSync can meet these requirements by providing fast and efficient data transfer with network throttling capability to control bandwidth usage.

upvoted 5 times

cloudbusting 2 years, 3 months ago

<https://docs.aws.amazon.com/datasync/latest/userguide/configure-bandwidth.html>

upvoted 3 times

bdp123 2 years, 3 months ago

Selected Answer: C

<https://aws.amazon.com/datasync/>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #302

Topic 1

A company wants to create a mobile app that allows users to stream slow-motion video clips on their mobile devices. Currently, the app captures video clips and uploads the video clips in raw format into an Amazon S3 bucket. The app retrieves these video clips directly from the S3 bucket. However, the videos are large in their raw format.


Users are experiencing issues with buffering and playback on mobile devices. The company wants to implement solutions to maximize the performance and scalability of the app while minimizing operational overhead.

Which combination of solutions will meet these requirements? (Choose two.)

- A. Deploy Amazon CloudFront for content delivery and caching. **Most Voted**
- B. Use AWS DataSync to replicate the video files across AWS Regions in other S3 buckets.
- C. Use Amazon Elastic Transcoder to convert the video files to more appropriate formats.
- D. Deploy an Auto Scaling group of Amazon EC2 instances in Local Zones for content delivery and caching.
- E. Deploy an Auto Scaling group of Amazon EC2 instances to convert the video files to more appropriate formats.

Correct Answer: A

Community vote distribution



Comments

pentium75 **Highly Voted** 1 year, 5 months ago

F - Fire the guy who created the current design
upvoted 22 times

awsgeek75 1 year, 4 months ago

No, make him watch all those videos with buffering!
upvoted 14 times

Bhawesh **Highly Voted** 2 years, 3 months ago

For Minimum operational overhead, the 2 options A,C should be correct.
A. Deploy Amazon CloudFront for content delivery and caching.

C. Use Amazon Elastic Transcoder to convert the video files to more appropriate formats.

upvoted 21 times

surajkrishnamurthy Most Recent 4 months, 1 week ago

Selected Answer: C

Not able to select Multi Option

Admin please check

upvoted 1 times

Rcosmos 4 months, 2 weeks ago

Selected Answer: U

A melhor combinação de soluções para atender aos requisitos é:

A. Implante o Amazon CloudFront para entrega de conteúdo e armazenamento em cache.

C. Use o Amazon Elastic Transcoder para converter os arquivos de vídeo em formatos mais apropriados.

upvoted 1 times

JA2018 6 months, 3 weeks ago

Selected Answer: C

Which combination of solutions will meet these requirements? (Choose two.)

.....

upvoted 1 times

Mayank0502 11 months, 2 weeks ago

Selected Answer: C

A & C. Admin has almost every answer wrong

upvoted 3 times

xyGGXH 1 year, 3 months ago

Selected Answer: A

A&C is correct

upvoted 4 times

db95476 1 year, 5 months ago

Selected Answer: A

A and C

upvoted 3 times

Ruffyit 1 year, 6 months ago

For Minimum operational overhead, the 2 options A,C should be correct.

A. Deploy Amazon CloudFront for content delivery and caching.

C. Use Amazon Elastic Transcoder to convert the video files to more appropriate formats.

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

For Minimum operational overhead, the 2 options A,C should be correct.

A. Deploy Amazon CloudFront for content delivery and caching.

C. Use Amazon Elastic Transcoder to convert the video files to more appropriate formats.

upvoted 2 times

Guru4Cloud 1 year, 8 months ago

examtopics team, please fix this question, please allow to select two answer

upvoted 3 times

vincent2023 4 months, 3 weeks ago

It has been more than 1year and 3 months, but the issue hasn't been fixed.

upvoted 1 times

upvoted 1 times

jacob_ho 1 year, 9 months ago

Elastic Transcoder has been deprecated, and AWS encourage to use AWS Elemental MediaConvert right now:
<https://aws.amazon.com/blogs/media/how-to-migrate-workflows-from-amazon-elastic-transcoder-to-aws-elemental-mediaconvert/>

upvoted 7 times

enc_0343 1 year, 11 months ago

Selected Answer: A

AC is the correct answer

upvoted 2 times

antropaws 2 years ago

Selected Answer: A

AC, the only possible answers.

upvoted 2 times

Eden 2 years, 1 month ago

It says choose two so I chose AC

upvoted 2 times

WheracanIstart 2 years, 2 months ago

Selected Answer: C

A & C are the right answers.

upvoted 5 times

kampatra 2 years, 2 months ago

Selected Answer: A

Correct answer: AC

upvoted 4 times

Steve_4542636 2 years, 3 months ago

Selected Answer: C

A and C. Transcoder does exactly what this needs.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #303

Topic 1

A company is launching a new application deployed on an Amazon Elastic Container Service (Amazon ECS) cluster and is using the Fargate launch type for ECS tasks. The company is monitoring CPU and memory usage because it is expecting high traffic to the application upon its launch. However, the company wants to reduce costs when utilization decreases.

What should a solutions architect recommend?

- A. Use Amazon EC2 Auto Scaling to scale at certain periods based on previous traffic patterns.
- B. Use an AWS Lambda function to scale Amazon ECS based on metric breaches that trigger an Amazon CloudWatch alarm.
- C. Use Amazon EC2 Auto Scaling with simple scaling policies to scale when ECS metric breaches trigger an Amazon CloudWatch alarm.
- D. Use AWS Application Auto Scaling with target tracking policies to scale when ECS metric breaches trigger an Amazon CloudWatch alarm. **Most Voted**

Correct Answer: D*Community vote distribution*

D (100%)

Comments**rrharris** **Highly Voted** 1 year, 9 months ago

Answer is D - Auto-scaling with target tracking
upvoted 14 times

Joxtat **Highly Voted** 1 year, 9 months ago**Selected Answer: D**

<https://docs.aws.amazon.com/autoscaling/application/userguide/what-is-application-auto-scaling.html>
upvoted 5 times

phuonglai **Most Recent** 10 months ago**Selected Answer: D**

<https://docs.aws.amazon.com/autoscaling/application/userguide/what-is-application-auto-scaling.html>
upvoted 4 times

awsgeek75 10 months, 3 weeks ago

Selected Answer: D

<https://docs.aws.amazon.com/autoscaling/application/userguide/what-is-application-auto-scaling.html>

upvoted 2 times

pentium75 11 months, 2 weeks ago

Selected Answer: D

This is running on Fargate, so EC2 scaling (A and C) is out. Lambda (B) is too complex.

upvoted 4 times

TariqKipkemei 1 year, 2 months ago

Target tracking will scale in/out the ECS cluster to maintain the average CPU utilization to a set value. e.g. <<50%>>> Scale out when average CPU utilization is above 50% until average CPU utilization is back to 50%. And scale in when average CPU utilization is below 50% until average CPU utilization is back to 50%.

upvoted 5 times

TariqKipkemei 1 year, 2 months ago

Answer is D

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

Answer is D - Auto-scaling with target tracking

upvoted 2 times

vincent2023 4 months, 3 weeks ago

Explanation please !!!

upvoted 1 times

TariqKipkemei 1 year, 7 months ago

Answer is D - Application Auto Scaling is a web service for developers and system administrators who need a solution for automatically scaling their scalable resources for individual AWS services beyond Amazon EC2.

upvoted 4 times

boxu03 1 year, 9 months ago

Selected Answer: D

should be D

upvoted 2 times

jahmad0730 1 year, 9 months ago

Selected Answer: D

Answer is D

upvoted 3 times

Neha999 1 year, 9 months ago

D : auto-scaling with target tracking

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #304

Topic 1

A company recently created a disaster recovery site in a different AWS Region. The company needs to transfer large amounts of data back and forth between NFS file systems in the two Regions on a periodic basis.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use AWS DataSync. **Most Voted**
- B. Use AWS Snowball devices.
- C. Set up an SFTP server on Amazon EC2.
- D. Use AWS Database Migration Service (AWS DMS).

Correct Answer: A

Community vote distribution

A (100%)

Comments

LuckyAro **Highly Voted** 1 year, 3 months ago

Selected Answer: A

AWS DataSync is a fully managed data transfer service that simplifies moving large amounts of data between on-premises storage systems and AWS services. It can also transfer data between different AWS services, including different AWS Regions. DataSync provides a simple, scalable, and automated solution to transfer data, and it minimizes the operational overhead because it is fully managed by AWS.

upvoted 18 times

Ruffyt **Most Recent** 6 months, 3 weeks ago

AWS DataSync is a fully managed data transfer service that simplifies moving large amounts of data between on-premises storage systems and AWS services. It can also transfer data between different AWS services, including different AWS Regions. DataSync provides a simple, scalable, and automated solution to transfer data, and it minimizes the operational overhead because it is fully managed by AWS.

upvoted 2 times

Salilgen 5 months, 3 weeks ago

You cannot use DataSync to transfer across Regions with an NFS

<https://docs.aws.amazon.com/datasync/latest/userguide/working-with-locations.html#working-with-locations-cross-regions>

upvoted 1 times

TariqKipkemei 8 months ago

Selected Answer: A

Use AWS DataSync

upvoted 2 times

Guru4Cloud 9 months ago

Selected Answer: A

Use AWS DataSync.

upvoted 2 times

kruasan 1 year, 1 month ago

Selected Answer: A

- AWS DataSync is a data transfer service optimized for moving large amounts of data between NFS file systems. It can automatically copy files and metadata between your NFS file systems in different AWS Regions.
- DataSync requires minimal setup and management. You deploy a source and destination agent, provide the source and destination locations, and DataSync handles the actual data transfer efficiently in the background.
- DataSync can schedule and monitor data transfers to keep source and destination in sync with minimal overhead. It resumes interrupted transfers and validates data integrity.
- DataSync optimizes data transfer performance across AWS's network infrastructure. It can achieve high throughput with minimal impact to your operations.

upvoted 4 times

kruasan 1 year, 1 month ago

Option B - AWS Snowball requires physical devices to transfer data. This incurs overhead to transport devices and manually load/unload data. It is not an online data transfer solution.

Option C - Setting up and managing an SFTP server would require provisioning EC2 instances, handling security groups, and writing scripts to automate the data transfer - all of which demand more overhead than DataSync.

Option D - AWS Database Migration Service is designed for migrating databases, not general file system data. It would require converting your NFS data into a database format, incurring additional overhead.

upvoted 4 times

ashu089 1 year, 2 months ago

Selected Answer: A

A only

upvoted 2 times

skiwili 1 year, 3 months ago

Selected Answer: A

Aaaaaa

upvoted 2 times

NolaHOla 1 year, 3 months ago

A should be correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #305

Topic 1

A company is designing a shared storage solution for a gaming application that is hosted in the AWS Cloud. The company needs the ability to use SMB clients to access data. The solution must be fully managed.

Which AWS solution meets these requirements?

- A. Create an AWS DataSync task that shares the data as a mountable file system. Mount the file system to the application server.
- B. Create an Amazon EC2 Windows instance. Install and configure a Windows file share role on the instance. Connect the application server to the file share.
- C. Create an Amazon FSx for Windows File Server file system. Attach the file system to the origin server. Connect the application server to the file system. **Most Voted**
- D. Create an Amazon S3 bucket. Assign an IAM role to the application to grant access to the S3 bucket. Mount the S3 bucket to the application server.

Correct Answer: C

Community vote distribution

C (100%)

Comments

rrharris **Highly Voted** 1 year, 9 months ago

Answer is C - SMB = storage gateway or FSx
upvoted 9 times

Neha999 **Highly Voted** 1 year, 9 months ago

C L: Amazon FSx for Windows File Server file system
upvoted 7 times

8883b6c **Most Recent** 5 months, 1 week ago

Selected Answer: C

SMB -> FSx
upvoted 1 times

phuonglai 10 months ago

Selected Answer: C

SMB -> FSx
upvoted 4 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: C

SMB = FSx for Windows File Server
upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

Answer is C - SMB = storage gateway or FSx
upvoted 3 times

kruasan 1 year, 7 months ago

Selected Answer: C

- Amazon FSx for Windows File Server provides a fully managed native Windows file system that can be accessed using the industry-standard SMB protocol. This allows Windows clients like the gaming application to directly access file data.
- FSx for Windows File Server handles time-consuming file system administration tasks like provisioning, setup, maintenance, file share management, backups, security, and software patching - reducing operational overhead.
- FSx for Windows File Server supports high file system throughput, IOPS, and consistent low latencies required for performance-sensitive workloads. This makes it suitable for a gaming application.
- The file system can be directly attached to EC2 instances, providing a performant shared storage solution for the gaming servers.

upvoted 5 times

kruasan 1 year, 7 months ago

Option A - DataSync is for data transfer, not providing a shared file system. It cannot be mounted or directly accessed.
Option B - A self-managed EC2 file share would require manually installing, configuring and maintaining a Windows file system and share. This demands significant overhead to operate.
Option D - Amazon S3 is object storage, not a native file system. The data in S3 would need to be converted/formatted to provide file share access, adding complexity. S3 cannot be directly mounted or provide the performance of FSx.

upvoted 5 times

elearningtakai 1 year, 8 months ago

Selected Answer: C

Amazon FSx for Windows File Server
upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: C

I vote C since FSx supports SMB
upvoted 2 times

LuckyAro 1 year, 9 months ago

Selected Answer: C

AWS FSx for Windows File Server is a fully managed native Microsoft Windows file system that is accessible through the SMB protocol. It provides features such as file system backups, integrated with Amazon S3, and Active Directory integration for user authentication and access control. This solution allows for the use of SMB clients to access the data and is fully managed, eliminating the need for the company to manage the underlying infrastructure.

upvoted 3 times

Babba 1 year, 9 months ago

Selected Answer: C

C for me
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #306

Topic 1

A company wants to run an in-memory database for a latency-sensitive application that runs on Amazon EC2 instances. The application processes more than 100,000 transactions each minute and requires high network throughput. A solutions architect needs to provide a cost-effective network design that minimizes data transfer charges.

Which solution meets these requirements?

- A. Launch all EC2 instances in the same Availability Zone within the same AWS Region. Specify a placement group with cluster strategy when launching EC2 instances. **Most Voted**
- B. Launch all EC2 instances in different Availability Zones within the same AWS Region. Specify a placement group with partition strategy when launching EC2 instances.
- C. Deploy an Auto Scaling group to launch EC2 instances in different Availability Zones based on a network utilization target.
- D. Deploy an Auto Scaling group with a step scaling policy to launch EC2 instances in different Availability Zones.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**kruasan** **Highly Voted** 1 year, 7 months ago**Selected Answer: A**

Reasons:

- Launching instances within a single AZ and using a cluster placement group provides the lowest network latency and highest bandwidth between instances. This maximizes performance for an in-memory database and high-throughput application.
- Communications between instances in the same AZ and placement group are free, minimizing data transfer charges. Inter-AZ and public IP traffic can incur charges.
- A cluster placement group enables the instances to be placed close together within the AZ, allowing the high network throughput required. Partition groups span AZs, reducing bandwidth.
- Auto Scaling across zones could launch instances in AZs that increase data transfer charges. It may reduce network throughput, impacting performance.

upvoted 23 times

kruasan 1 year, 7 months ago

In contrast:

Option B - A partition placement group spans AZs, reducing network bandwidth between instances and potentially increasing costs.

Option C - Auto Scaling alone does not guarantee the network throughput and cost controls required for this use case.

Launching across AZs could increase data transfer charges.

Option D - Step scaling policies determine how many instances to launch based on metrics alone. They lack control over network connectivity and costs between instances after launch.

upvoted 12 times

TariqKipkemei Highly Voted 1 year, 2 months ago

Selected Answer: A

Cluster placement group packs instances close together inside an Availability Zone. This strategy enables workloads to achieve the low-latency network performance.

upvoted 5 times

awsgeek75 Most Recent 10 months, 3 weeks ago

Selected Answer: A

Apart from the fact that BCD distribute the instances across AZ which is bad for inter-node network latency, I think the following article is really useful in understanding A:

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>

upvoted 2 times

Ruffyit 1 year ago

- Launching instances within a single AZ and using a cluster placement group provides the lowest network latency and highest bandwidth between instances. This maximizes performance for an in-memory database and high-throughput application.
- Communications between instances in the same AZ and placement group are free, minimizing data transfer charges. Inter-AZ and public IP traffic can incur charges.
- A cluster placement group enables the instances to be placed close together within the AZ, allowing the high network throughput required. Partition groups span AZs, reducing bandwidth.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

Launch all EC2 instances in the same Availability Zone within the same AWS Region. Specify a placement group with cluster strategy when launching EC2 instances

upvoted 2 times

NoinNothing 1 year, 7 months ago

Selected Answer: A

Cluster - have low latency if its in same AZ and same region so Answer is "A"

upvoted 3 times

BeeKayEnn 1 year, 8 months ago

Answer would be A - As part of selecting all the EC2 instances in the same availability zone, they all will be within the same DC and logically the latency will be very less as compared to the other Availability Zones..

As all the autoscaling nodes will also be on the same availability zones, (as per Placement groups with Cluster mode), this would provide the low-latency network performance

Reference is below:

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>

upvoted 4 times

[Removed] 1 year, 8 months ago

Selected Answer: A

A - Low latency, high net throughput

upvoted 2 times

elearningtakai 1 year, 8 months ago

Selected Answer: A

A placement group is a logical grouping of instances within a single Availability Zone, and it provides low-latency network connectivity between instances. By launching all EC2 instances in the same Availability Zone and specifying a placement group with cluster strategy, the application can take advantage of the high network throughput and low latency network connectivity that placement groups provide.

upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: A

Cluster placement groups improves throughput between the instances which means less EC2 instances would be needed thus reducing costs.

upvoted 2 times

maciekmaciek 1 year, 9 months ago

Selected Answer: A

A because Specify a placement group

upvoted 2 times

KZM 1 year, 9 months ago

It is option A:

To achieve low latency, high throughput, and cost-effectiveness, the optimal solution is to launch EC2 instances as a placement group with the cluster strategy within the same Availability Zone.

upvoted 3 times

ManOnTheMoon 1 year, 9 months ago

Why not C?

upvoted 2 times

Steve_4542636 1 year, 9 months ago

You're thinking operational efficiency. The question asks for cost reduction.

upvoted 4 times

rrharris 1 year, 9 months ago

Answer is A - Clustering

upvoted 3 times

Neha999 1 year, 9 months ago

A : Cluster placement group

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #307

Topic 1

A company that primarily runs its application servers on premises has decided to migrate to AWS. The company wants to minimize its need to scale its Internet Small Computer Systems Interface (iSCSI) storage on premises. The company wants only its recently accessed data to remain stored locally.

Which AWS solution should the company use to meet these requirements?

- A. Amazon S3 File Gateway
- B. AWS Storage Gateway Tape Gateway
- C. AWS Storage Gateway Volume Gateway stored volumes
- D. AWS Storage Gateway Volume Gateway cached volumes **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

LuckyAro **Highly Voted** 2 years, 3 months ago

Selected Answer: D

AWS Storage Gateway Volume Gateway provides two configurations for connecting to iSCSI storage, namely, stored volumes and cached volumes. The stored volume configuration stores the entire data set on-premises and asynchronously backs up the data to AWS. The cached volume configuration stores recently accessed data on-premises, and the remaining data is stored in Amazon S3.

Since the company wants only its recently accessed data to remain stored locally, the cached volume configuration would be the most appropriate. It allows the company to keep frequently accessed data on-premises and reduce the need for scaling its iSCSI storage while still providing access to all data through the AWS cloud. This configuration also provides low-latency access to frequently accessed data and cost-effective off-site backups for less frequently accessed data.

upvoted 41 times

tonybuivannghia 8 months ago

You're correct but a small mistake that the stored volume configuration stores the entire data set on-premises and asynchronously backs up the data to AWS (S3). The cached volume is low latency access to most recent data.

upvoted 2 times

smgsi Highly Voted 2 years, 3 months ago

Selected Answer: D

https://docs.amazonaws.cn/en_us/storagegateway/latest/vgw/StorageGatewayConcepts.html#storage-gateway-cached-concepts

upvoted 10 times

TariqKipkemei Most Recent 1 year, 8 months ago

Selected Answer: D

Frequently accessed data = AWS Storage Gateway Volume Gateway cached volumes

upvoted 4 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

The best AWS solution to meet the requirements is to use AWS Storage Gateway cached volumes (option D).

The key points:

Company migrating on-prem app servers to AWS

Want to minimize scaling on-prem iSCSI storage

Only recent data should remain on-premises

The AWS Storage Gateway cached volumes allow the company to connect their on-premises iSCSI storage to AWS cloud storage. It stores frequently accessed data locally in the cache for low-latency access, while older data is stored in AWS.

upvoted 4 times

kruasan 2 years, 1 month ago

Selected Answer: D

- Volume Gateway cached volumes store entire datasets on S3, while keeping a portion of recently accessed data on your local storage as a cache. This meets the goal of minimizing on-premises storage needs while keeping hot data local.
- The cache provides low-latency access to your frequently accessed data, while long-term retention of the entire dataset is provided durable and cost-effective in S3.
- You get virtually unlimited storage on S3 for your infrequently accessed data, while controlling the amount of local storage used for cache. This simplifies on-premises storage scaling.
- Volume Gateway cached volumes support iSCSI connections from on-premises application servers, allowing a seamless migration experience. Servers access local cache and S3 storage volumes as iSCSI LUNs.

upvoted 8 times

kruasan 2 years, 1 month ago

In contrast:

Option A - S3 File Gateway only provides file interfaces (NFS/SMB) to data in S3. It does not support block storage or cache recently accessed data locally.

Option B - Tape Gateway is designed for long-term backup and archiving to virtual tape cartridges on S3. It does not provide primary storage volumes or local cache for low-latency access.

Option C - Volume Gateway stored volumes keep entire datasets locally, then asynchronously back them up to S3. This does not meet the goal of minimizing on-premises storage needs.

upvoted 6 times

Steve_4542636 2 years, 3 months ago

Selected Answer: D

I vote D

upvoted 2 times

ManOnTheMoon 2 years, 3 months ago

Agree with D

upvoted 2 times

Babba 2 years, 3 months ago

Selected Answer: D

recently accessed data to remain stored locally - cached

upvoted 4 times

Bhawesh 2 years, 3 months ago

Selected Answer: D

D. AWS Storage Gateway Volume Gateway cached volumes
upvoted 4 times

bdp123 2 years, 3 months ago

Selected Answer: D

recently accessed data to remain stored locally - cached
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #308

Topic 1

A company has multiple AWS accounts that use consolidated billing. The company runs several active high performance Amazon RDS for Oracle On-Demand DB instances for 90 days. The company's finance team has access to AWS Trusted Advisor in the consolidated billing account and all other AWS accounts.

The finance team needs to use the appropriate AWS account to access the Trusted Advisor check recommendations for RDS. The finance team must review the appropriate Trusted Advisor check to reduce RDS costs.

Which combination of steps should the finance team take to meet these requirements? (Choose two.)

- A. Use the Trusted Advisor recommendations from the account where the RDS instances are running.
- B. Use the Trusted Advisor recommendations from the consolidated billing account to see all RDS instance checks at the same time. **Most Voted**
- C. Review the Trusted Advisor check for Amazon RDS Reserved Instance Optimization.
- D. Review the Trusted Advisor check for Amazon RDS Idle DB Instances. **Most Voted**
- E. Review the Trusted Advisor check for Amazon Redshift Reserved Node Optimization.

Correct Answer: BD

Community vote distribution



Comments

Nietzsche82 **Highly Voted** 2 years, 3 months ago

Selected Answer: BD

B & D

<https://aws.amazon.com/premiumsupport/knowledge-center/trusted-advisor-cost-optimization/>

upvoted 19 times

Michal_L_95 **Highly Voted** 2 years, 3 months ago

Selected Answer: BC

I would go with B and C as the company is running for 90 days and C option is basing on 30 days report which would mean that there is higher potential on cost saving rather than on idle instances

upvoted 6 times

Steve_4542636 2 years, 3 months ago

C is stating "Reserved Instances" The question states they are using On Demand Instances. Reserved instances are reserved for less money for 1 or 3 years.

upvoted 8 times

Lalo 2 years ago

In the scenario it says for 90 days, therefore the correct answer is D

No C

upvoted 2 times

pentium75 1 year, 5 months ago

D checks for instances that did not have a connection for a while, which is not the case here, all instances are active.

upvoted 2 times

Michal_L_95 2 years, 2 months ago

Once read the question again, I agree with you.

upvoted 1 times

pentium75 1 year, 5 months ago

C is about optimizing on-demand instances by turning them into reserved instances.

upvoted 2 times

c12ab95 Most Recent 2 weeks, 5 days ago

Selected Answer: AC

A. Use the Trusted Advisor recommendations from the account where the RDS instances are running. Trusted Advisor cost optimization checks for RDS (such as Reserved Instance Optimization and Idle DB Instances) are available in the account where the RDS resources are provisioned. The finance team must access the recommendations from the relevant account to see the checks for those specific RDS instances.

C. Review the Trusted Advisor check for Amazon RDS Reserved Instance Optimization. This check analyzes your On-Demand RDS usage and recommends purchasing Reserved Instances to save costs. It is a key Trusted Advisor recommendation for reducing RDS costs, especially for workloads running for extended periods

upvoted 1 times

Yak_Yeti 1 month, 1 week ago

Selected Answer: AC

The article emphasizes that Trusted Advisor generates recommendations at the individual AWS account level. While the consolidated billing account provides a centralized view of billing, the specific optimization recommendations for resources like RDS instances will be found within the account where those resources reside.

The article also mentions checks for idle resources like "Amazon EC2 Idle Instances" and "Amazon RDS Idle DB Instances." However, as the original analysis pointed out, the scenario specifies that the RDS instances are active and running a high-performance application, making the "Idle DB Instances" check less relevant for the immediate goal of reducing costs for these specific instances.

upvoted 1 times

a8a1e0e 4 months, 1 week ago

Selected Answer: AC

A: AWS Trusted Advisor recommendations for RDS cost optimization are only available within the AWS account where the instances are running. The consolidated billing (payer) account does not provide detailed service-level checks for linked accounts. Therefore, the finance team must access the Trusted Advisor check directly from the AWS account that owns the RDS instances to see relevant recommendations.

C: The company has been using high-performance RDS instances for 90 days, which indicates a long-term workload. Reserved Instances (RIs) offer significant cost savings compared to On-Demand pricing when committing to one or three years. Trusted Advisor identifies instances with consistent usage and recommends RIs to reduce costs, making this the most relevant check for optimizing expenses. Even after 90 days, purchasing an RI for the remaining duration still provides substantial savings.

upvoted 5 times

MatAlves 9 months ago

Selected Answer: BC

The answer is either BC or BD, depending on how you interpret "The company runs several active... instances for 90 days for 90 days."

D: it assumes the instances will only run for 90 days, so reserved instances can't be the answer, since it requires 1-3 years utilization.

C: it assumes there is no idle instances since they've been active for the last 90 days.

upvoted 4 times

aagarwallko 9 months ago

B: Use the Trusted Advisor recommendations from the consolidated billing account to see all RDS instance checks at the same time. This option allows the finance team to see all RDS instance checks across all AWS accounts in one place. Since the company uses consolidated billing, this account will have access to all of the AWS accounts' Trusted Advisor recommendations.

C: Review the Trusted Advisor check for Amazon RDS Reserved Instance Optimization. This check can help identify cost savings opportunities for RDS by identifying instances that can be covered by Reserved Instances. This can result in significant savings on RDS costs.

upvoted 2 times

ike001 11 months, 2 weeks ago

BD is the answer-Amazon Redshift Reserved Node Optimization and Relational Database Service (RDS) Reserved Instance Optimization check are not available to accounts linked in consolidated billing.

upvoted 3 times

sandordini 1 year, 1 month ago

Selected Answer: BD

"you can reserve a DB instance for a one- or three-year term". We only have data for 90 days. I feel it too risky to commit for 1/3 year(s) without information on future usage. If we knew that we expected the same usage pattern for the next 1,2,3 years, I'd agree with C.

upvoted 3 times

soufiyane 1 year, 1 month ago

Selected Answer: BC

B) Use the Trusted Advisor recommendations from the consolidated billing account to see all RDS instance checks at the same time. This option allows the finance team to see all RDS instance checks across all AWS accounts in one place. Since the company uses consolidated billing, this account will have access to all of the AWS accounts' Trusted Advisor recommendations.

C) Review the Trusted Advisor check for Amazon RDS Reserved Instance Optimization. This check can help identify cost savings opportunities for RDS by identifying instances that can be covered by Reserved Instances. This can result in significant savings on RDS costs.

upvoted 2 times

Rhydian25 11 months, 3 weeks ago

Reserved Instances are for 1 or 3 years. Not for 90 days

upvoted 5 times

scar0909 1 year, 3 months ago

Selected Answer: BC

<https://aws.amazon.com/premiumsupport/knowledge-center/trusted-advisor-cost-optimization/>

upvoted 2 times

bujuman 1 year, 3 months ago

Selected Answer: BD

Insights: The company runs several active high performance Amazon RDS for Oracle On-Demand DB instances for 90 days So it's clear that this company need to check the configuration of any Amazon Relational Database Service (Amazon RDS) for any database (DB) instances that appear to be idle.

upvoted 2 times

dkw2342 1 year, 3 months ago

0 0 0

BC

AWS Trusted Advisor is an online resource to help you reduce cost, increase performance, and improve security by optimizing your AWS environment. (...) Recommendations are based on the previous calendar month's hour-by-hour usage aggregated across all consolidated billing accounts.

<https://docs.aws.amazon.com/whitepapers/latest/cost-optimization-reservation-models/aws-trusted-advisor.html>

Amazon EC2 Reserved Instance Optimization: An important part of using AWS involves balancing your Reserved Instance (RI) purchase against your On-Demand Instance usage. This check provides recommendations on which RIs will help reduce the costs incurred from using On-Demand Instances. We create these recommendations by analyzing your On-Demand usage for the past 30 days. We then categorizing the usage into eligible categories for reservations.

<https://docs.aws.amazon.com/awssupport/latest/user/cost-optimization-checks.html#amazon-ec2-reserved-instances-optimization>

upvoted 2 times

Salilgen 5 months, 3 weeks ago

The company run instances only for 90 days then reserved instances cannot save costs

upvoted 1 times

NayeraB 1 year, 3 months ago

Selected Answer: BC

If you're choosing D for the idle instances, Amazon RDS Reserved Instance Optimization Trusted Advisor check includes recommendations related to underutilized and idle RDS instances. It helps identify instances that are not fully utilized and provides recommendations on how to optimize costs, such as resizing or terminating unused instances, or purchasing reserved instances to match usage patterns more efficiently.

upvoted 3 times

leejwli 1 year, 4 months ago

Selected Answer: BC

Reserved Instances can be shared across accounts, and that is the reason why we need to check the consolidated bill.

upvoted 3 times

farnamjam 1 year, 5 months ago

Selected Answer: BC

BC

we don't want to check Idle instances because the instances were active for last 90 days. Idle means it was inactive for at least 7 days.

upvoted 5 times

farnamjam 1 year, 5 months ago

Selected Answer: BD

BD

we don't want to check Idle instances because the instances were active for last 90 days. Idle means it was inactive for at least 7 days.

upvoted 2 times

MatAlves 9 months ago

Then why you voted for D?

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #309

Topic 1


A solutions architect needs to optimize storage costs. The solutions architect must identify any Amazon S3 buckets that are no longer being accessed or are rarely accessed.

Which solution will accomplish this goal with the LEAST operational overhead?

- A. Analyze bucket access patterns by using the S3 Storage Lens dashboard for advanced activity metrics. **Most Voted**
- B. Analyze bucket access patterns by using the S3 dashboard in the AWS Management Console.
- C. Turn on the Amazon CloudWatch BucketSizeBytes metric for buckets. Analyze bucket access patterns by using the metrics data with Amazon Athena.
- D. Turn on AWS CloudTrail for S3 object monitoring. Analyze bucket access patterns by using CloudTrail logs that are integrated with Amazon CloudWatch Logs.

Correct Answer: A

Community vote distribution



Comments

kpato87 **Highly Voted** 2 years, 3 months ago

Selected Answer: A

S3 Storage Lens is a fully managed S3 storage analytics solution that provides a comprehensive view of object storage usage, activity trends, and recommendations to optimize costs. Storage Lens allows you to analyze object access patterns across all of your S3 buckets and generate detailed metrics and reports.

upvoted 23 times

bdp123 **Highly Voted** 2 years, 3 months ago

Selected Answer: A

<https://aws.amazon.com/blogs/aws/s3-storage-lens/>

upvoted 5 times

Gape4 **Most Recent** 11 months, 2 weeks ago

Selected Answer: A

S3 Storage Lens includes an interactive dashboard which you can find in the S3 console. The dashboard gives you the ability to

perform filtering and drill-down into your metrics to really understand how your storage is being used. The metrics are organized into categories like data protection and cost efficiency, to allow you to easily find relevant metrics.

upvoted 2 times

AmijoSando 1 year ago

Anyone passed the exam can confirm the right answer ? A or D

upvoted 1 times

xyGGXH 1 year, 3 months ago

Selected Answer: A

A
S3 Storage Lens is the first cloud storage analytics solution to provide a single view of object storage usage and activity across hundreds, or even thousands, of accounts in an organization, with drill-downs to generate insights at multiple aggregation levels.

upvoted 3 times

Neung983 1 year, 3 months ago

On the other hand, Option B suggests using the S3 dashboard in the AWS Management Console, which provides a straightforward and user-friendly interface to monitor S3 bucket access patterns. This option may have less operational overhead compared to setting up and managing Storage Lens. Additionally, for simply identifying rarely accessed buckets, the built-in metrics and access analysis provided by the S3 dashboard can often suffice without the need for advanced analytics offered by Storage Lens. Therefore, Option B is considered to have less operational overhead for the specific task described in the question.

upvoted 1 times

jaswantn 1 year, 4 months ago

But nowhere on S3 Storage Lens dashboard this information is available; that when the bucket is accessed last time. But it gives insight on the bucket's size. with this information we can check if files can be moved to less costly storage class. This way we can reduce storage cost..... The information which is the main requirement of the given scenario, is available when we use Cloudtrail logs ... so i choose option D.

upvoted 1 times

jaswantn 1 year, 4 months ago

if the bucket is being accessed frequently then we can leave it as it is, otherwise we can move the files to infrequent access storage class thus can save some money.

upvoted 2 times

Ruffyit 1 year, 6 months ago

S3 Storage Lens is a fully managed S3 storage analytics solution that provides a comprehensive view of object storage usage, activity trends, and recommendations to optimize costs. Storage Lens allows you to analyze object access patterns across all of your S3 buckets and generate detailed metrics and reports.

upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: A

Amazon S3 Storage Lens was designed to handle this requirement.

upvoted 2 times

Wayne23Fang 1 year, 9 months ago

Selected Answer: D

A missed turning on monitoring. It can also help you learn about your customer base and understand your Amazon S3 bill. By default, Amazon S3 doesn't collect server access logs. When you enable logging, Amazon S3 delivers access logs for a source bucket to a target bucket that you choose.

I could not find that S3 storage Lens examples online showing using Lens to identify idle S3 buckets. Instead I find using S3 Access Logging. Hmm.

upvoted 3 times

pentium75 1 year, 5 months ago

How will you find when a bucket was used the last time if you turn on logging NOW?

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

S3 Storage Lens is a cloud-storage analytics feature that provides you with 29+ usage and activity metrics, including object count, size, age, and access patterns. This data can help you understand how your data is being used and identify areas where you can optimize your storage costs.

The S3 Storage Lens dashboard provides an interactive view of your storage usage and activity trends. This makes it easy to identify buckets that are no longer being accessed or are rarely accessed.

The S3 Storage Lens dashboard is a fully managed service, so there is no need to set up or manage any additional infrastructure.

upvoted 3 times

BigHammer 1 year, 9 months ago

"S3 Storage Lens" seems to be the popular answer, however, where in Storage Lens can you see if a bucket/object is being USED? I see all kinds of stats, but not that.

upvoted 3 times

MatAlves 9 months ago

"S3 Storage Lens delivers organization-wide visibility into object storage usage, activity trends, and makes actionable recommendations to optimize costs..."

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

<https://aws.amazon.com/blogs/aws/s3-storage-lens/>

upvoted 3 times

kruasan 2 years, 1 month ago

Selected Answer: A

The S3 Storage Lens dashboard provides visibility into storage metrics and activity patterns to help optimize storage costs. It shows metrics like objects added, objects deleted, storage consumed, and requests. It can filter by bucket, prefix, and tag to analyze specific subsets of data

upvoted 4 times

kruasan 2 years, 1 month ago

B) The standard S3 console dashboard provides basic info but would require manually analyzing metrics for each bucket. This does not scale well and requires significant overhead.

C) Turning on the BucketSizeBytes metric and analyzing the data in Athena may provide insights but would require enabling metrics, building Athena queries, and analyzing the results. This requires more operational effort than option A.

D) Enabling CloudTrail logging and monitoring the logs in CloudWatch Logs could provide access pattern data but would require setting up CloudTrail, monitoring the logs, and analyzing the relevant info. This option has the highest operational overhead

upvoted 5 times

LuckyAro 2 years, 3 months ago

Selected Answer: A

S3 Storage Lens provides a dashboard with advanced activity metrics that enable the identification of infrequently accessed and unused buckets. This can help a solutions architect optimize storage costs without incurring additional operational overhead.

upvoted 4 times

Babba 2 years, 3 months ago

Selected Answer: A

it looks like it's A

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #310

Topic 1

A company sells datasets to customers who do research in artificial intelligence and machine learning (AI/ML). The datasets are large, formatted files that are stored in an Amazon S3 bucket in the us-east-1 Region. The company hosts a web application that the customers use to purchase access to a given dataset. The web application is deployed on multiple Amazon EC2 instances behind an Application Load Balancer. After a purchase is made, customers receive an S3 signed URL that allows access to the files.

The customers are distributed across North America and Europe. The company wants to reduce the cost that is associated with data transfers and wants to maintain or improve performance.

What should a solutions architect do to meet these requirements?

- A. Configure S3 Transfer Acceleration on the existing S3 bucket. Direct customer requests to the S3 Transfer Acceleration endpoint. Continue to use S3 signed URLs for access control.
- B. Deploy an Amazon CloudFront distribution with the existing S3 bucket as the origin. Direct customer requests to the CloudFront URL. Switch to CloudFront signed URLs for access control. **Most Voted**
- C. Set up a second S3 bucket in the eu-central-1 Region with S3 Cross-Region Replication between the buckets. Direct customer requests to the closest Region. Continue to use S3 signed URLs for access control.
- D. Modify the web application to enable streaming of the datasets to end users. Configure the web application to read the data from the existing S3 bucket. Implement access control directly in the application.

Correct Answer: B

Community vote distribution

B (100%)

Comments

LuckyAro **Highly Voted** 1 year, 9 months ago

Selected Answer: B

To reduce the cost associated with data transfers and maintain or improve performance, a solutions architect should use Amazon CloudFront, a content delivery network (CDN) service that securely delivers data, videos, applications, and APIs to customers globally with low latency and high transfer speeds.

Deploying a CloudFront distribution with the existing S3 bucket as the origin will allow the company to serve the data to

Deploying a CloudFront distribution with the existing S3 bucket as the origin will allow the company to serve the data to customers from edge locations that are closer to them, reducing data transfer costs and improving performance.

Directing customer requests to the CloudFront URL and switching to CloudFront signed URLs for access control will enable customers to access the data securely and efficiently.

upvoted 14 times

awsgeek75 Highly Voted 10 months, 3 weeks ago

Selected Answer: B

A: Speeds uploads

C: Increases the cost rather than reducing it

D: Stopped reading after "Modify the web application..."

upvoted 10 times

Ruffyt Most Recent 1 year ago

To reduce the cost associated with data transfers and maintain or improve performance, a solutions architect should use Amazon CloudFront, a content delivery network (CDN) service that securely delivers data, videos, applications, and APIs to customers globally with low latency and high transfer speeds.

upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: B

Technically both option B and C will work. But because cost is a factor then Amazon CloudFront should be the preferred option.

upvoted 2 times

react97 1 year, 2 months ago

Selected Answer: B

- B.
1. Amazon CloudFront caches content at edge locations -- reducing the need for frequent data transfer from S3 bucket -- thus significantly lowering data transfer costs (as compared to directly serving data from S3 bucket to customers in different regions)
 2. CloudFront delivers content to users from the nearest edge location -- minimizing latency -- improves performance for customers

A - focus on accelerating uploads to S3 which may not necessarily improve the performance needed for serving datasets to customers

C - helps with redundancy and data availability but does not necessarily offer cost savings for data transfer.

D - complex to implement, does not address data transfer cost

upvoted 6 times

bdp123 1 year, 9 months ago

Selected Answer: B

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/PrivateContent.html>

upvoted 4 times

Bhawesh 1 year, 9 months ago

Selected Answer: B

B. Deploy an Amazon CloudFront distribution with the existing S3 bucket as the origin. Direct customer requests to the CloudFront URL. Switch to CloudFront signed URLs for access control.

<https://www.examttopics.com/discussions/amazon/view/68990-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #311

Topic 1

A company is using AWS to design a web application that will process insurance quotes. Users will request quotes from the application. Quotes must be separated by quote type, must be responded to within 24 hours, and must not get lost. The solution must maximize operational efficiency and must minimize maintenance.

Which solution meets these requirements?

- A. Create multiple Amazon Kinesis data streams based on the quote type. Configure the web application to send messages to the proper data stream. Configure each backend group of application servers to use the Kinesis Client Library (KCL) to pool messages from its own data stream.
- B. Create an AWS Lambda function and an Amazon Simple Notification Service (Amazon SNS) topic for each quote type. Subscribe the Lambda function to its associated SNS topic. Configure the application to publish requests for quotes to the appropriate SNS topic.
- C. Create a single Amazon Simple Notification Service (Amazon SNS) topic. Subscribe Amazon Simple Queue Service (Amazon SQS) queues to the SNS topic. Configure SNS message filtering to publish messages to the proper SQS queue based on the quote type. Configure each backend application server to use its own SQS queue. **Most Voted**
- D. Create multiple Amazon Kinesis Data Firehose delivery streams based on the quote type to deliver data streams to an Amazon OpenSearch Service cluster. Configure the application to send messages to the proper delivery stream. Configure each backend group of application servers to search for the messages from OpenSearch Service and process them accordingly.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**LuckyAro** **Highly Voted** 2 years, 3 months ago**Selected Answer: C**

Quote types need to be separated: SNS message filtering can be used to publish messages to the appropriate SQS queue based on the quote type, ensuring that quotes are separated by type. Quotes must be responded to within 24 hours and must not get lost: SQS provides reliable and scalable queuing for messages, ensuring that quotes will not get lost and can be processed in a timely manner. Additionally, each backend application server can use its own SQS queue, ensuring that quotes are processed efficiently without any delay.

can use its own SQS queue, ensuring that quotes are processed efficiently without any delay.

Operational efficiency and minimizing maintenance: Using a single SNS topic and multiple SQS queues is a scalable and cost-effective approach, which can help to maximize operational efficiency and minimize maintenance. Additionally, SNS and SQS are fully managed services, which means that the company will not need to worry about maintenance tasks such as software updates, hardware upgrades, or scaling the infrastructure.

upvoted 19 times

Vlad Highly Voted 2 years, 3 months ago

C is the best option

upvoted 9 times

akshay243007 Most Recent 10 months, 2 weeks ago

Selected Answer: C

SQS + SNS = fanout

upvoted 3 times

Uzbekistan 1 year, 3 months ago

Option C would be the most suitable solution to meet the requirements while maximizing operational efficiency and minimizing maintenance.

Explanation:

Amazon SNS (Simple Notification Service) allows for the creation of a single topic to which multiple subscribers can be attached. In this scenario, each quote type can be considered a subscriber. Amazon SQS (Simple Queue Service) queues can be subscribed to the SNS topic, and SNS message filtering can be used to direct messages to the appropriate SQS queue based on the quote type. This setup ensures that quotes are separated by quote type and that they are not lost. Each backend application server can then poll its own SQS queue to retrieve and process messages. This architecture is efficient, scalable, and requires minimal maintenance, as it leverages managed AWS services without the need for complex custom code or infrastructure setup.

upvoted 4 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

I originally went for D due to searching requirements but Open Search is for analytics and logs and nothing to do with data coming from streams as in this question.

upvoted 2 times

Ruffyt 1 year, 6 months ago

Quote types need to be separated: SNS message filtering can be used to publish messages to the appropriate SQS queue based on the quote type, ensuring that quotes are separated by type.

Quotes must be responded to within 24 hours and must not get lost: SQS provides reliable and scalable queuing for messages, ensuring that quotes will not get lost and can be processed in a timely manner. Additionally, each backend application server can use its own SQS queue, ensuring that quotes are processed efficiently without any delay.

Operational efficiency and minimizing maintenance: Using a single SNS topic and multiple SQS queues is a scalable and cost-effective approach, which can help to maximize operational efficiency and minimize maintenance. Additionally, SNS and SQS are fully managed services, which means that the company will not need to worry about maintenance tasks such as software updates, hardware upgrades, or scaling the

upvoted 3 times

tekjm 1 year, 8 months ago

Keyword is "...and must not get lost" = SQS

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

Create a single SNS topic

Subscribe separate SQS queues per quote type

Use SNS message filtering to send messages to proper queue

Backend servers poll their respective SQS queue

The key points:

Quote requests must be processed within 24 hrs without loss

Need to maximize efficiency and minimize maintenance

Requests separated by quote type

upvoted 2 times

lexotan 2 years, 1 month ago

Selected Answer: C

This wrong answers from examtopic are getting me so frustrated. Which one is the correct answer then?

upvoted 6 times

Steve_4542636 2 years, 3 months ago

Selected Answer: C

This is the SNS fan-out technique where you will have one SNS service to many SQS services

<https://docs.aws.amazon.com/sns/latest/dg/sns-sqs-as-subscriber.html>

upvoted 7 times

UnluckyDucky 2 years, 2 months ago

SNS Fan-out fans message to all subscribers, this uses SNS filtering to publish the message only to the right SQS queue (not all of them).

upvoted 3 times

Yechi 2 years, 3 months ago

Selected Answer: C

<https://aws.amazon.com/getting-started/hands-on/filter-messages-published-to-topics/>

upvoted 8 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #312

Topic 1

A company has an application that runs on several Amazon EC2 instances. Each EC2 instance has multiple Amazon Elastic Block Store (Amazon EBS) data volumes attached to it. The application's EC2 instance configuration and data need to be backed up nightly. The application also needs to be recoverable in a different AWS Region.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. Write an AWS Lambda function that schedules nightly snapshots of the application's EBS volumes and copies the snapshots to a different Region.
- B. Create a backup plan by using AWS Backup to perform nightly backups. Copy the backups to another Region. Add the application's EC2 instances as resources. **Most Voted**
- C. Create a backup plan by using AWS Backup to perform nightly backups. Copy the backups to another Region. Add the application's EBS volumes as resources.
- D. Write an AWS Lambda function that schedules nightly snapshots of the application's EBS volumes and copies the snapshots to a different Availability Zone.

Correct Answer: B*Community vote distribution*

B (94%)

C (6%)

Comments**hasport** **Highly Voted** 1 year, 9 months ago

B is answer so the requirement is "The application's EC2 instance configuration and data need to be backed up nightly" so we need "add the application's EC2 instances as resources". This option will backup both EC2 configuration and data
upvoted 21 times

TungPham **Highly Voted** 1 year, 9 months ago**Selected Answer: B**

<https://aws.amazon.com/vi/blogs/aws/aws-backup-ec2-instances-efs-single-file-restore-and-cross-region-backup/>
When you back up an EC2 instance, AWS Backup will protect all EBS volumes attached to the instance, and it will attach them to an AMI that stores all parameters from the original EC2 instance except for two
upvoted 16 times

ravmondfekrv **Most Recent** 11 months, 3 weeks ago

Selected Answer: B

Question says: " The application's EC2 instance configuration and data need to be backed up", thus C is not correct, B is upvoted 3 times

Ruffyit 1 year ago

<https://aws.amazon.com/vi/blogs/aws/aws-backup-ec2-instances-efs-single-file-restore-and-cross-region-backup/>
When you back up an EC2 instance, AWS Backup will protect all EBS volumes attached to the instance, and it will attach them to an AMI that stores all parameters from the original EC2 instance except for two
upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: B

As part of configuring a backup plan you need to enable (opt-in) resource types that will be protected by the backup plan. For this case EC2.

<https://aws.amazon.com/getting-started/hands-on/amazon-ec2-backup-and-restore-using-aws-backup/#:~:text=the%20services%20used%20with-,AWS%20Backup,-a.%20In%20the%20navigation>
upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

B is the most appropriate solution because it allows you to create a backup plan to automate the backup process of EC2 instances and EBS volumes, and copy backups to another region. Additionally, you can add the application's EC2 instances as resources to ensure their configuration and data are backed up nightly.

upvoted 2 times

Geekboii 1 year, 8 months ago

i would say B
upvoted 2 times

Geekboii 1 year, 8 months ago

i would say B
upvoted 2 times

AlmeroSenior 1 year, 9 months ago

Selected Answer: B

AWS KB states if you select the EC2 instance , associated EBS's will be auto covered .

<https://aws.amazon.com/blogs/aws/aws-backup-ec2-instances-efs-single-file-restore-and-cross-region-backup/>
upvoted 3 times

LuckyAro 1 year, 9 months ago

Selected Answer: B

B is the most appropriate solution because it allows you to create a backup plan to automate the backup process of EC2 instances and EBS volumes, and copy backups to another region. Additionally, you can add the application's EC2 instances as resources to ensure their configuration and data are backed up nightly.
A and D involve writing custom Lambda functions to automate the snapshot process, which can be complex and require more maintenance effort. Moreover, these options do not provide an integrated solution for managing backups and recovery, and copying snapshots to another region.

Option C involves creating a backup plan with AWS Backup to perform backups for EBS volumes only. This approach would not back up the EC2 instances and their configuration

upvoted 3 times

Mia2009687 1 year, 5 months ago

The data is stored in the EBS storage volume, EC2 won't hold the data, I think we need "Add the application's EBS volumes as resources."

upvoted 2 times

everfly 1 year, 9 months ago

Selected Answer: C

The application's EC2 instance configuration and data are stored on EBS volume right?

upvoted 2 times

pentium75 11 months, 2 weeks ago

No, this is not how EC2 works.

upvoted 1 times

awsgeek75 10 months, 3 weeks ago

No, ECS config is the config you provide when launching the EC2 instance. EBS is a resource for EC2 as a part of configuration. When you backup EC2, it will backup the instance which resulted from the configuration and that will include the EBS volumes that are attached to the instance.

upvoted 3 times

Rehan33 1 year, 9 months ago

The data is store on EBS volume so why we are not using EBS as a source instead of EC2

upvoted 1 times

thewalker 10 months, 3 weeks ago

Also, if EBS volumes are added or removed as the requirement, not need to update the AWS Config.

upvoted 1 times

obatunde 1 year, 9 months ago

Because "The application's EC2 instance configuration and data need to be backed up nightly"

upvoted 6 times

fulingyu288 1 year, 9 months ago

Selected Answer: B

Use AWS Backup to create a backup plan that includes the EC2 instances, Amazon EBS snapshots, and any other resources needed for recovery. The backup plan can be configured to run on a nightly schedule.

upvoted 2 times

zTopic 1 year, 9 months ago

Selected Answer: B

The application's EC2 instance configuration and data need to be backed up nightly >> B

upvoted 2 times

NolaHOla 1 year, 9 months ago

But isn't the data needed to be backed up on the EBS ?

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #313

Topic 1

A company is building a mobile app on AWS. The company wants to expand its reach to millions of users. The company needs to build a platform so that authorized users can watch the company's content on their mobile devices.

What should a solutions architect recommend to meet these requirements?

- A. Publish content to a public Amazon S3 bucket. Use AWS Key Management Service (AWS KMS) keys to stream content.
- B. Set up IPsec VPN between the mobile app and the AWS environment to stream content.
- C. Use Amazon CloudFront. Provide signed URLs to stream content. **Most Voted**
- D. Set up AWS Client VPN between the mobile app and the AWS environment to stream content.

Correct Answer: C

Community vote distribution



Comments

Steve_4542636 **Highly Voted** 1 year, 9 months ago

Selected Answer: C

Enough with CloudFront already.
upvoted 29 times

awsgeek75 10 months, 3 weeks ago

This whole exam seems like a sales pitch for CloudFront and SQS... lol!
upvoted 7 times

TariqKipkemei 1 year, 7 months ago

Hahaha..cloudfront too hyped :)
upvoted 3 times

LuckyAro **Highly Voted** 1 year, 9 months ago

Selected Answer: C

Amazon CloudFront is a content delivery network (CDN) that securely delivers data, videos, applications, and APIs to customers globally with low latency and high transfer speeds. CloudFront supports signed URLs that provide authorized access to your

greatest network latency and high transfer speeds. CloudFront supports signed URLs that provide authorized access to your content. This feature allows the company to control who can access their content and for how long, providing a secure and scalable solution for millions of users.

upvoted 9 times

mwwt2022 11 months, 1 week ago

great explanation!

upvoted 2 times

lostmagnet001 Most Recent 10 months ago

Selected Answer: C

CF always for reaching places

upvoted 3 times

Ruffyt 1 year ago

Use Amazon CloudFront. Provide signed URLs to stream content.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

Use Amazon CloudFront. Provide signed URLs to stream content.

upvoted 2 times

antropaws 1 year, 6 months ago

Selected Answer: C

C is correct.

upvoted 3 times

kprakashbehera 1 year, 9 months ago

Cloudfront is the correct solution.

upvoted 5 times

datz 1 year, 8 months ago

Feel your pain :D hahaha

upvoted 2 times

jennyka76 1 year, 9 months ago

C

<https://www.amazonaws.cn/en/cloudfront/>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #314

Topic 1

A company has an on-premises MySQL database used by the global sales team with infrequent access patterns. The sales team requires the database to have minimal downtime. A database administrator wants to migrate this database to AWS without selecting a particular instance type in anticipation of more users in the future.

Which service should a solutions architect recommend?

- A. Amazon Aurora MySQL
- B. Amazon Aurora Serverless for MySQL **Most Voted**
- C. Amazon Redshift Spectrum
- D. Amazon RDS for MySQL

Correct Answer: B

Community vote distribution

B (100%)

Comments

cloudbusting **Highly Voted** 1 year, 9 months ago

"without selecting a particular instance type" = serverless
upvoted 30 times

elearningtakai **Highly Voted** 1 year, 8 months ago

Selected Answer: B

With Aurora Serverless for MySQL, you don't need to select a particular instance type, as the service automatically scales up or down based on the application's needs.
upvoted 11 times

awsgeek75 **Most Recent** 10 months, 3 weeks ago

Selected Answer: B

The DBA had one job and he doesn't want to do it... so B it is
upvoted 9 times

Ruffyt 1 year ago

without selecting a particular instance type = Amazon Aurora Serverless for MySQL
upvoted 4 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: B

without selecting a particular instance type = Amazon Aurora Serverless for MySQL
upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

B. Amazon Aurora Serverless for MySQL
upvoted 2 times

Diqian 1 year, 3 months ago

What's the difference between A and B. I think Aurora is serverless, isn't it?
upvoted 1 times

Valder21 1 year, 3 months ago

seems serverless is an option of amazon aurora. Not a very good naming scheme.
upvoted 2 times

Srikanth0057 1 year, 9 months ago

Selected Answer: B

Bbbbbbb
upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: B

<https://aws.amazon.com/rds/aurora/serverless/>
upvoted 2 times

LuckyAro 1 year, 9 months ago

Selected Answer: B

Amazon Aurora Serverless for MySQL is a fully managed, auto-scaling relational database service that scales up or down automatically based on the application demand. This service provides all the capabilities of Amazon Aurora, such as high availability, durability, and security, without requiring the customer to provision any database instances.

With Amazon Aurora Serverless for MySQL, the sales team can enjoy minimal downtime since the database is designed to automatically scale to accommodate the increased traffic. Additionally, the service allows the customer to pay only for the capacity used, making it cost-effective for infrequent access patterns.

Amazon RDS for MySQL could also be an option, but it requires the customer to select an instance type, and the database administrator would need to monitor and adjust the instance size manually to accommodate the increasing traffic.
upvoted 3 times

Drayen25 1 year, 9 months ago

Minimal downtime points directly to Aurora Serverless
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #315

Topic 1

A company experienced a breach that affected several applications in its on-premises data center. The attacker took advantage of vulnerabilities in the custom applications that were running on the servers. The company is now migrating its applications to run on Amazon EC2 instances. The company wants to implement a solution that actively scans for vulnerabilities on the EC2 instances and sends a report that details the findings.

Which solution will meet these requirements?

- A. Deploy AWS Shield to scan the EC2 instances for vulnerabilities. Create an AWS Lambda function to log any findings to AWS CloudTrail.
- B. Deploy Amazon Macie and AWS Lambda functions to scan the EC2 instances for vulnerabilities. Log any findings to AWS CloudTrail.
- C. Turn on Amazon GuardDuty. Deploy the GuardDuty agents to the EC2 instances. Configure an AWS Lambda function to automate the generation and distribution of reports that detail the findings.
- D. Turn on Amazon Inspector. Deploy the Amazon Inspector agent to the EC2 instances. Configure an AWS Lambda function to automate the generation and distribution of reports that detail the findings. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****siyam008** **Highly Voted** 1 year, 9 months ago**Selected Answer: D**

AWS Shield for DDOS

Amazon Macie for discover and protect sensitive data

Amazon GuardDuty for intelligent threat discovery to protect AWS account

Amazon Inspector for automated security assessment. like known Vulnerability

upvoted 58 times

robotgeek 4 months, 2 weeks ago

So 58 upvotes and you are talking about "known Vulnerability" for a "custom app"... ok

upvoted 1 times

upvoted 1 times

benacert Highly Voted 12 months ago

Whenever I feel vulnerable, I use AWS Inspector..

upvoted 17 times

zinabu Most Recent 7 months, 2 weeks ago

Selected Answer: D

Amazon Inspector for automated security assessment. like known Vulnerability

upvoted 4 times

Ruffyt 1 year ago

AWS Shield for DDOS

Amazon Macie for discover and protect sensitive data

Amazon GuardDuty for intelligent threat discovery to protect AWS account

Amazon Inspector for automated security assessment. like known Vulnerability

upvoted 5 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: D

vulnerabilities = Amazon Inspector

malicious activity = Amazon GuardDuty

upvoted 11 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

Enable Amazon Inspector

Deploy Inspector agents to EC2 instances

Use Lambda to generate and distribute vulnerability reports

The key points:

Migrate on-prem apps with vulnerabilities to EC2

Need active scanning of EC2 instances for vulnerabilities

Require reports on findings

upvoted 4 times

kruasan 1 year, 7 months ago

Selected Answer: D

Amazon Inspector:

- Performs active vulnerability scans of EC2 instances. It looks for software vulnerabilities, unintended network accessibility, and other security issues.
- Requires installing an agent on EC2 instances to perform scans. The agent must be deployed to each instance.
- Provides scheduled scan reports detailing any findings of security risks or vulnerabilities. These reports can be used to patch or remediate issues.
- Is best suited for proactively detecting security weaknesses and misconfigurations in your AWS environment.

upvoted 4 times

kruasan 1 year, 7 months ago

Amazon GuardDuty:

- Monitors for malicious activity like unusual API calls, unauthorized infrastructure deployments, or compromised EC2 instances. It uses machine learning and behavioral analysis of logs.
- Does not require installing any agents. It relies on analyzing AWS CloudTrail, VPC Flow Logs, and DNS logs.
- Alerts you to any detected threats, suspicious activity or policy violations in your AWS accounts. These alerts warrant investigation but may not always require remediation.
- Is focused on detecting active threats, unauthorized behavior, and signs of a compromise in your AWS environment.
- Can also detect some vulnerabilities and misconfigurations but coverage is not as broad as a dedicated service like Inspector.

upvoted 7 times

datz 1 year, 8 months ago

Selected Answer: D

Amazon Inspector is a vulnerability scanning tool that you can use to identify potential security issues within your EC2 instances.

It is a kind of automated security assessment service that checks the network exposure of your EC2 or latest security state for applications running into your EC2 instance. It has ability to auto discover your AWS workload and continuously scan for the open loophole or vulnerability.

upvoted 2 times

shanwford 1 year, 8 months ago

Selected Answer: D

Amazon Inspector is a vulnerability scanning tool that you can use to identify potential security issues within your EC2 instances. Guard Duty continuously monitors your entire AWS account via Cloud Trail, Flow Logs, DNS Logs as Input.

upvoted 2 times

GalileoEC2 1 year, 8 months ago

Selected Answer: C

:) C is the correct

<https://cloudkatha.com/amazon-guardduty-vs-inspector-which-one-should-you-use/>

upvoted 1 times

jayantp04 11 months, 3 weeks ago

Document itself saying that
Amazon Inspector is a vulnerability scanning tool
hence correct Answer is D

upvoted 2 times

MssP 1 year, 8 months ago

Please, read the link you sent: Amazon Inspector is a vulnerability scanning tool that you can use to identify potential security issues within your EC2 instances. GuardDuty is very critical part to identify threats, based on that findings you can setup automated preventive actions or remediation's. So Answer is D.

upvoted 2 times

GalileoEC2 1 year, 8 months ago

Selected Answer: C

<https://cloudkatha.com/amazon-guardduty-vs-inspector-which-one-should-you-use/>

upvoted 1 times

LuckyAro 1 year, 9 months ago

Selected Answer: D

Amazon Inspector is a security assessment service that helps to identify security vulnerabilities and compliance issues in applications deployed on Amazon EC2 instances. It can be used to assess the security of applications that are deployed on Amazon EC2 instances, including those that are custom-built.

To use Amazon Inspector, the Amazon Inspector agent must be installed on the EC2 instances that need to be assessed. The agent collects data about the instances and sends it to Amazon Inspector for analysis. Amazon Inspector then generates a report that details any security vulnerabilities that were found and provides guidance on how to remediate them.

By configuring an AWS Lambda function, the company can automate the generation and distribution of reports that detail the findings. This means that reports can be generated and distributed as soon as vulnerabilities are detected, allowing the company to take action quickly.

upvoted 2 times

pbpally 1 year, 9 months ago

Selected Answer: D

I'm a little confused on how someone came up with C, it is definitely D.

upvoted 2 times

obatunde 1 year, 9 months ago

Selected Answer: D

Amazon Inspector

upvoted 3 times

obatunde 1 year, 9 months ago

Amazon Inspector is an automated vulnerability management service that continually scans AWS workloads for software vulnerabilities and unintended network exposure. <https://aws.amazon.com/inspector/features/?nc=sn&loc=2>
upvoted 4 times

Palanda 1 year, 9 months ago

Selected Answer: D

I think D
upvoted 1 times

minglu 1 year, 9 months ago

Selected Answer: D

Inspector for EC2
upvoted 1 times

skiwili 1 year, 9 months ago

Selected Answer: D

Ddddddd
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #316

Topic 1


A company uses an Amazon EC2 instance to run a script to poll for and process messages in an Amazon Simple Queue Service (Amazon SQS) queue. The company wants to reduce operational costs while maintaining its ability to process a growing number of messages that are added to the queue.

What should a solutions architect recommend to meet these requirements?

- A. Increase the size of the EC2 instance to process messages faster.
- B. Use Amazon EventBridge to turn off the EC2 instance when the instance is underutilized.
- C. Migrate the script on the EC2 instance to an AWS Lambda function with the appropriate runtime. **Most Voted**
- D. Use AWS Systems Manager Run Command to run the script on demand.

Correct Answer: C

Community vote distribution



Comments

kpato87 **Highly Voted** 2 years, 3 months ago

Selected Answer: C

By migrating the script to AWS Lambda, the company can take advantage of the auto-scaling feature of the service. AWS Lambda will automatically scale resources to match the size of the workload. This means that the company will not have to worry about provisioning or managing instances as the number of messages increases, resulting in lower operational costs

upvoted 13 times

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: C

The key points are:

Currently using an EC2 instance to poll SQS and process messages

Want to reduce costs while handling growing message volume

By migrating the polling script to a Lambda function, the company can avoid the cost of running a dedicated EC2 instance.

Lambda functions scale automatically to handle message spikes. And Lambda billing is based on actual usage, resulting in cost savings versus provisioned EC2 capacity.

upvoted 8 times

TariqKipkemei Most Recent 1 year, 8 months ago

Selected Answer: C

reduce operational costs = serverless = Lambda functions
upvoted 2 times

Steve_4542636 2 years, 3 months ago

Selected Answer: C

Lambda costs money only when it's processing, not when idle
upvoted 3 times

ManOnTheMoon 2 years, 3 months ago

Agree with C
upvoted 1 times

khasport 2 years, 3 months ago

the answer is C. With this option, you can reduce operational cost as the question mentioned
upvoted 1 times

LuckyAro 2 years, 3 months ago

Selected Answer: C

AWS Lambda is a serverless compute service that allows you to run your code without provisioning or managing servers. By migrating the script to an AWS Lambda function, you can eliminate the need to maintain an EC2 instance, reducing operational costs. Additionally, Lambda automatically scales to handle the increasing number of messages in the SQS queue.
upvoted 1 times

zTopic 2 years, 3 months ago

Selected Answer: C

It Should be C.
Lambda allows you to execute code without provisioning or managing servers, so it is ideal for running scripts that poll for and process messages in an Amazon SQS queue. The scaling of the Lambda function is automatic, and you only pay for the actual time it takes to process the messages.
upvoted 3 times

Bhawesh 2 years, 3 months ago

Selected Answer: D

To reduce the operational overhead, it should be:
D. Use AWS Systems Manager Run Command to run the script on demand.
upvoted 3 times

lucdt4 2 years ago

No, replace EC2 instead by using lambda to reduce costs
upvoted 1 times

pentium75 1 year, 5 months ago

So every time an item is added to the queue, you log into AWS Systems Manager through your browser, select "Run Command" and select your instance and enter the command to run the script?
upvoted 4 times

ike001 11 months, 2 weeks ago

very sarcastic question :)
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #317

Topic 1

A company uses a legacy application to produce data in CSV format. The legacy application stores the output data in Amazon S3. The company is deploying a new commercial off-the-shelf (COTS) application that can perform complex SQL queries to analyze data that is stored in Amazon Redshift and Amazon S3 only. However, the COTS application cannot process the .csv files that the legacy application produces.

The company cannot update the legacy application to produce data in another format. The company needs to implement a solution so that the COTS application can use the data that the legacy application produces.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an AWS Glue extract, transform, and load (ETL) job that runs on a schedule. Configure the ETL job to process the .csv files and store the processed data in Amazon Redshift. **Most Voted**

B. Develop a Python script that runs on Amazon EC2 instances to convert the .csv files to .sql files. Invoke the Python script on a cron schedule to store the output files in Amazon S3.

C. Create an AWS Lambda function and an Amazon DynamoDB table. Use an S3 event to invoke the Lambda function. Configure the Lambda function to perform an extract, transform, and load (ETL) job to process the .csv files and store the processed data in the DynamoDB table.

D. Use Amazon EventBridge to launch an Amazon EMR cluster on a weekly schedule. Configure the EMR cluster to perform an extract, transform, and load (ETL) job to process the .csv files and store the processed data in an Amazon Redshift table.

Correct Answer: A*Community vote distribution*

A (98%)

C

Comments**awsgeek75** **Highly Voted** 10 months, 3 weeks ago**Selected Answer: A**

Time to sell some Glue.

I believe these kind of questions are there to indoctrinate us into acknowledging how blessed we are to have managed services

like AWS Glue when you look at other horrible and painful options

upvoted 19 times

elearningtakai Highly Voted 1 year, 8 months ago

Selected Answer: A

A, AWS Glue is a fully managed ETL service that can extract data from various sources, transform it into the required format, and load it into a target data store. In this case, the ETL job can be configured to read the CSV files from Amazon S3, transform the data into a format that can be loaded into Amazon Redshift, and load it into an Amazon Redshift table.

B requires the development of a custom script to convert the CSV files to SQL files, which could be time-consuming and introduce additional operational overhead. C, while using serverless technology, requires the additional use of DynamoDB to store the processed data, which may not be necessary if the data is only needed in Amazon Redshift. D, while an option, is not the most efficient solution as it requires the creation of an EMR cluster, which can be costly and complex to manage.

upvoted 8 times

pentium75 Most Recent 11 months, 2 weeks ago

Selected Answer: A

B - Developing a script is surely not minimizing operational effort

C - Stores data in DynamoDB where the new app cannot use it

D - Could work but is total overkill (EMR is for Big Data analysis, not for simple ETL)

upvoted 4 times

Ruffyt 1 year ago

A-ETL is serverless & best suited with the requirement who primary job is ETL

B-Usage of Ec2 adds operational overhead & incur costs

C-DynamoDB(NoSql) does suit the requirement as company is performing SQL queries

D-EMR adds operational overhead & incur costs

upvoted 3 times

ACloud_Guru15 1 year, 1 month ago

Selected Answer: A

A-ETL is serverless & best suited with the requirement who primary job is ETL

B-Usage of Ec2 adds operational overhead & incur costs

C-DynamoDB(NoSql) does suit the requirement as company is performing SQL queries

D-EMR adds operational overhead & incur costs

upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: A

Data transformation = AWS Glue

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

Create an AWS Glue ETL job to process the CSV files

Configure the job to run on a schedule

Output the transformed data to Amazon Redshift

The key points:

Legacy app generates CSV files in S3

New app requires data in Redshift or S3

Need to transform CSV to support new app with minimal ops overhead

upvoted 2 times

kraken21 1 year, 8 months ago

Selected Answer: A

Glue is server less and has less operational head than EMR so A.

upvoted 2 times

[Removed] 1 year, 8 months ago

Selected Answer: C

o meet the requirement with the least operational overhead. a serverless approach should be used. Among the options

to meet the requirements with the least operational overhead, a serverless approach should be used. Among the options provided, option C provides a serverless solution using AWS Lambda, S3, and DynamoDB. Therefore, the solution should be to create an AWS Lambda function and an Amazon DynamoDB table. Use an S3 event to invoke the Lambda function. Configure the Lambda function to perform an extract, transform, and load (ETL) job to process the .csv files and store the processed data in the DynamoDB table.

Option A is also a valid solution, but it may involve more operational overhead than Option C. With Option A, you would need to set up and manage an AWS Glue job, which would require more setup time than creating an AWS Lambda function. Additionally, AWS Glue jobs have a minimum execution time of 10 minutes, which may not be necessary or desirable for this use case. However, if the data processing is particularly complex or requires a lot of data transformation, AWS Glue may be a more appropriate solution.

upvoted 1 times

pentium75 11 months, 2 weeks ago

Creating and maintaining a Lambda function is more "operational overhead" than using a ready-made service such as Glue. But more important, answer C says "store the processed data in the DynamoDB table" while the application can "analyze data that is stored in Amazon Redshift and Amazon S3 only".

upvoted 2 times

MssP 1 year, 8 months ago

Important point: The COTS performs complex SQL queries to analyze data in Amazon Redshift. If you use DynamoDB -> No SQL queries. Option A makes more sense.

upvoted 4 times

LuckyAro 1 year, 9 months ago

Selected Answer: A

A would be the best solution as it involves the least operational overhead. With this solution, an AWS Glue ETL job is created to process the .csv files and store the processed data directly in Amazon Redshift. This is a serverless approach that does not require any infrastructure to be provisioned, configured, or maintained. AWS Glue provides a fully managed, pay-as-you-go ETL service that can be easily configured to process data from S3 and load it into Amazon Redshift. This approach allows the legacy application to continue to produce data in the CSV format that it currently uses, while providing the new COTS application with the ability to analyze the data using complex SQL queries.

upvoted 4 times

jennyka76 1 year, 9 months ago

A

<https://docs.aws.amazon.com/glue/latest/dg/aws-glue-programming-etl-format-csv-home.html>

I AGREE AFTER READING LINK

upvoted 2 times

cloudbusting 1 year, 9 months ago

A: <https://docs.aws.amazon.com/glue/latest/dg/aws-glue-programming-etl-format.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #318

Topic 1


A company recently migrated its entire IT environment to the AWS Cloud. The company discovers that users are provisioning oversized Amazon EC2 instances and modifying security group rules without using the appropriate change control process. A solutions architect must devise a strategy to track and audit these inventory and configuration changes.

Which actions should the solutions architect take to meet these requirements? (Choose two.)

- A. Enable AWS CloudTrail and use it for auditing. **Most Voted**
- B. Use data lifecycle policies for the Amazon EC2 instances.
- C. Enable AWS Trusted Advisor and reference the security dashboard.
- D. Enable AWS Config and create rules for auditing and compliance purposes. **Most Voted**
- E. Restore previous resource configurations with an AWS CloudFormation template.

Correct Answer: AD

Community vote distribution



Comments

LuckyAro **Highly Voted** 1 year, 9 months ago

Selected Answer: AD

A. Enable AWS CloudTrail and use it for auditing. CloudTrail provides event history of your AWS account activity, including actions taken through the AWS Management Console, AWS Command Line Interface (CLI), and AWS SDKs and APIs. By enabling CloudTrail, the company can track user activity and changes to AWS resources, and monitor compliance with internal policies and external regulations.

D. Enable AWS Config and create rules for auditing and compliance purposes. AWS Config provides a detailed inventory of the AWS resources in your account, and continuously records changes to the configurations of those resources. By creating rules in AWS Config, the company can automate the evaluation of resource configurations against desired state, and receive alerts when configurations drift from compliance.

Options B, C, and E are not directly relevant to the requirement of tracking and auditing inventory and configuration changes.
upvoted 14 times

zdi561 **Most Recent** 4 months ago

Selected Answer: CD

Amazon Web Services (AWS) CloudTrail is enabled by default for all AWS accounts.

upvoted 1 times

Ruffyt 1 year ago

A. Enable AWS CloudTrail and use it for auditing. CloudTrail provides event history of your AWS account activity, including actions taken through the AWS Management Console, AWS Command Line Interface (CLI), and AWS SDKs and APIs. By enabling CloudTrail, the company can track user activity and changes to AWS resources, and monitor compliance with internal policies and external regulations.

D. Enable AWS Config and create rules for auditing and compliance purposes. AWS Config provides a detailed inventory of the AWS resources in your account, and continuously records changes to the configurations of those resources. By creating rules in AWS Config, the company can automate the evaluation of resource configurations against desired state, and receive alerts when configurations drift from compliance.

Options B, C, and E are not directly relevant to the requirement of tracking and auditing inventory and configuration changes.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: AD

A. Enable AWS CloudTrail and use it for auditing. CloudTrail provides event history of your AWS account activity, including actions taken through the AWS Management Console, AWS Command Line Interface (CLI), and AWS SDKs and APIs. By enabling CloudTrail, the company can track user activity and changes to AWS resources, and monitor compliance with internal policies and external regulations.

D. Enable AWS Config and create rules for auditing and compliance purposes. AWS Config provides a detailed inventory of the AWS resources in your account, and continuously records changes to the configurations of those resources. By creating rules in AWS Config, the company can automate the evaluation of resource configurations against desired state, and receive alerts when configurations drift from compliance.

upvoted 2 times

mrsoa 1 year, 4 months ago

Selected Answer: CD

I am gonna go with CD

AWS Cloudtrail is already enabled so no need to enable it and for the auditing we are gonna use AWS config Answer D

C because Trusted advisor checks the security groups

upvoted 1 times

pentium75 11 months, 2 weeks ago

"AWS CloudTrail is already enabled" says who?

upvoted 3 times

awsgeek75 10 months, 3 weeks ago

CloudTrail is not enabled by default or in the question scenario. Even if it was, Trusted Advisor would just give you recommendations and usage reports. It won't audit anything for you

upvoted 1 times

kruasan 1 year, 7 months ago

Selected Answer: AD

A) Enable AWS CloudTrail and use it for auditing.

AWS CloudTrail provides a record of API calls and can be used to audit changes made to EC2 instances and security groups. By analyzing CloudTrail logs, the solutions architect can track who provisioned oversized instances or modified security groups without proper approval.

D) Enable AWS Config and create rules for auditing and compliance purposes.

AWS Config can record the configuration changes made to resources like EC2 instances and security groups. The solutions architect can create AWS Config rules to monitor for non-compliant changes, like launching certain instance types or opening security group ports without permission. AWS Config would alert on any violations of these rules.

upvoted 3 times

kruasan 1 year, 7 months ago

The other options would not fully meet the auditing and change tracking requirements:

B) Data lifecycle policies control when EC2 instances are backed up or deleted but do not audit configuration changes.

by Data being the primary concern when IAM instances are locked up or deleted but do not have configuration changes.

C) AWS Trusted Advisor security checks may detect some compliance violations after the fact but do not comprehensively log changes like AWS CloudTrail and AWS Config do.

E) CloudFormation templates enable rollback but do not provide an audit trail of changes. The solutions architect would not know who made unauthorized modifications in the first place.

upvoted 3 times

skiwili 1 year, 9 months ago

Selected Answer: AD

Yes A and D

upvoted 2 times

jennyka76 1 year, 9 months ago

AGREE WITH ANSWER - A & D
CloudTrail and Config

upvoted 2 times

Neha999 1 year, 9 months ago

CloudTrail and Config

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #319

Topic 1


A company has hundreds of Amazon EC2 Linux-based instances in the AWS Cloud. Systems administrators have used shared SSH keys to manage the instances. After a recent audit, the company's security team is mandating the removal of all shared keys. A solutions architect must design a solution that provides secure access to the EC2 instances.

Which solution will meet this requirement with the LEAST amount of administrative overhead?

- A. Use AWS Systems Manager Session Manager to connect to the EC2 instances. **Most Voted**
- B. Use AWS Security Token Service (AWS STS) to generate one-time SSH keys on demand.
- C. Allow shared SSH access to a set of bastion instances. Configure all other instances to allow only SSH access from the bastion instances.
- D. Use an Amazon Cognito custom authorizer to authenticate users. Invoke an AWS Lambda function to generate a temporary SSH key.

Correct Answer: A

Community vote distribution



Comments

Vlad **Highly Voted** 1 year, 9 months ago

Answer is A

Using AWS Systems Manager Session Manager to connect to the EC2 instances is a secure option as it eliminates the need for inbound SSH ports and removes the requirement to manage SSH keys manually. It also provides a complete audit trail of user activity. This solution requires no additional software to be installed on the EC2 instances.

upvoted 12 times

pentium75 **Highly Voted** 11 months, 2 weeks ago

Selected Answer: A

A - Systems Manager Session Manager has EXACTLY that purpose, 'providing secure access to EC2 instances'

B - STS can generate temporary IAM credentials or access keys but NOT SSH keys

C - Does not 'remove all shared keys' as requested

D - Cognito is not meant for internal users, and whole setup is complex

upvoted 7 times

pentium75 **Most Recent** 11 months, 2 weeks ago

pentium75 11 months, 2 weeks ago

Selected Answer: A

B - Querying is just a feature of Redshift but primarily it's a Data Warehouse - the question says nothing that historical data would have to be stored or accessed or analyzed

upvoted 2 times

Ruffyit 1 year ago

The key reasons why:

STS can generate short-lived credentials that provide temporary access to the EC2 instances for administering them.

The credentials can be generated on-demand each time access is needed, eliminating the risks of using permanent shared SSH keys.

No infrastructure like bastion hosts needs to be maintained.

The on-premises administrators can use the familiar SSH tools with the temporary keys.

upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: A

Session Manager provides secure and auditable node management without the need to open inbound ports, maintain bastion hosts, or manage SSH keys.

upvoted 2 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: B

The key reasons why:

STS can generate short-lived credentials that provide temporary access to the EC2 instances for administering them.

The credentials can be generated on-demand each time access is needed, eliminating the risks of using permanent shared SSH keys.

No infrastructure like bastion hosts needs to be maintained.

The on-premises administrators can use the familiar SSH tools with the temporary keys.

upvoted 1 times

pentium75 11 months, 2 weeks ago

STS provides temporary IAM credentials, not SSH keys

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Using AWS Security Token Service (AWS STS) to generate one-time SSH keys on demand is a secure and efficient way to provide access to the EC2 instances without the need for shared SSH keys. STS is a fully managed service that can be used to generate temporary security credentials, allowing systems administrators to connect to the EC2 instances without having to share SSH keys. The temporary credentials can be generated on demand, reducing the administrative overhead associated with managing SSH access

upvoted 2 times

ofinto 1 year, 2 months ago

Can you please provide documentation about generating a one-time SSH with STS?

upvoted 2 times

kruasan 1 year, 7 months ago

Selected Answer: A

AWS Systems Manager Session Manager provides secure shell access to EC2 instances without the need for SSH keys. It meets the security requirement to remove shared SSH keys while minimizing administrative overhead.

upvoted 2 times

Guru4Cloud 1 year, 2 months ago

If the systems administrators need to access the EC2 instances from an on-premises environment, using Session Manager may not be the ideal solution.

upvoted 1 times

upvoted 1 times

kruasan 1 year, 7 months ago

Session Manager is a fully managed AWS Systems Manager capability. With Session Manager, you can manage your Amazon Elastic Compute Cloud (Amazon EC2) instances, edge devices, on-premises servers, and virtual machines (VMs). You can use either an interactive one-click browser-based shell or the AWS Command Line Interface (AWS CLI). Session Manager provides secure and auditable node management without the need to open inbound ports, maintain bastion hosts, or manage SSH keys. Session Manager also allows you to comply with corporate policies that require controlled access to managed nodes, strict security practices, and fully auditable logs with node access details, while providing end users with simple one-click cross-platform access to your managed nodes.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html>

upvoted 3 times

kruasan 1 year, 7 months ago

Who should use Session Manager?

Any AWS customer who wants to improve their security and audit posture, reduce operational overhead by centralizing access control on managed nodes, and reduce inbound node access.

Information Security experts who want to monitor and track managed node access and activity, close down inbound ports on managed nodes, or allow connections to managed nodes that don't have a public IP address.

Administrators who want to grant and revoke access from a single location, and who want to provide one solution to users for Linux, macOS, and Windows Server managed nodes.

Users who want to connect to a managed node with just one click from the browser or AWS CLI without having to provide SSH keys.

upvoted 3 times

Stanislav4907 1 year, 8 months ago

Selected Answer: C

You guys seriously don't want to go to SMSM for Every Single EC2. You have to create solution not used services for one time access. Bastion will give you option to manage 1000s EC2 machines from 1. Plus you can use Ansible from it.

upvoted 2 times

UnluckyDucky 1 year, 8 months ago

Session Manager is the best practice and recommended way by Amazon to manage your instances. Bastion hosts require remote access therefore exposing them to the internet.

The most secure way is definitely session manager therefore answer A is correct imho.

upvoted 4 times

Zox42 1 year, 8 months ago

Question: "the company's security team is mandating the removal of all shared keys", answer C can't be right because it says: "Allow shared SSH access to a set of bastion instances".

upvoted 7 times

Steve_4542636 1 year, 9 months ago

Selected Answer: A

I vote a

upvoted 2 times

LuckyAro 1 year, 9 months ago

Selected Answer: A

AWS Systems Manager Session Manager provides secure and auditable instance management without the need for any inbound connections or open ports. It allows you to manage your instances through an interactive one-click browser-based shell or through the AWS CLI. This means that you don't have to manage any SSH keys, and you don't have to worry about securing access to your instances as access is controlled through IAM policies.

upvoted 5 times

bdp123 1 year, 9 months ago

Selected Answer: A

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html>

https://docs.aws.amazon.com/systems-manager/latest/userguide/session-management.html
upvoted 3 times

jahmad0730 1 year, 9 months ago

Selected Answer: A

Answer must be A
upvoted 3 times

jennyka76 1 year, 9 months ago

ANSWER - A
AWS SESSION MANAGER IS CORRECT LEAST EFFORTS TO ACCESS LINUX SYSTEM IN AWS CONSOLE AND YOUR ARE ALREADY LOGIN TO AWS. SO NO NEED FOR THE TOKEN OR OTHER STUFF DONE IN THE BACKGROUND BY AWS. MAKES SENSE.
upvoted 3 times

cloudbusting 1 year, 9 months ago

Answer is A
upvoted 4 times

zTopic 1 year, 9 months ago

Selected Answer: A

Answer is A
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #320

Topic 1

A company is using a fleet of Amazon EC2 instances to ingest data from on-premises data sources. The data is in JSON format and ingestion rates can be as high as 1 MB/s. When an EC2 instance is rebooted, the data in-flight is lost. The company's data science team wants to query ingested data in near-real time.

Which solution provides near-real-time data querying that is scalable with minimal data loss?

- A. Publish data to Amazon Kinesis Data Streams, Use Kinesis Data Analytics to query the data. **Most Voted**
- B. Publish data to Amazon Kinesis Data Firehose with Amazon Redshift as the destination. Use Amazon Redshift to query the data.
- C. Store ingested data in an EC2 instance store. Publish data to Amazon Kinesis Data Firehose with Amazon S3 as the destination. Use Amazon Athena to query the data.
- D. Store ingested data in an Amazon Elastic Block Store (Amazon EBS) volume. Publish data to Amazon ElastiCache for Redis. Subscribe to the Redis channel to query the data.

Correct Answer: A*Community vote distribution*

A (74%)

B (26%)

Comments**LuckyAro** **Highly Voted** 2 years, 3 months ago**Selected Answer: A**

A: is the solution for the company's requirements. Publishing data to Amazon Kinesis Data Streams can support ingestion rates as high as 1 MB/s and provide real-time data processing. Kinesis Data Analytics can query the ingested data in real-time with low latency, and the solution can scale as needed to accommodate increases in ingestion rates or querying needs. This solution also ensures minimal data loss in the event of an EC2 instance reboot since Kinesis Data Streams has a persistent data store for up to 7 days by default.

upvoted 15 times

bogobob **Highly Voted** 1 year, 6 months ago**Selected Answer: B**

The fact they specifically mention "near real-time" twice tells me the correct answer is KDF. On top of which its easier to setup and maintain. KDS is really only needed if you need real-time. Also using redshift will mean permanent data retention. The data in A could be lost after a year. Redshift queries are slow but you're still querying near real-time data

It could be lost after a year. Redshift queries are slow but you're still querying near real-time data
upvoted 6 times

Ernestokoro 1 year, 6 months ago

You are very correct. see supporting link <https://jayendrapatil.com/aws-kinesis-data-streams-vs-kinesis-firehose/#:~:text=vs%20Kine...,Purpose,into%20AWS%20products%20for%20processing>.

upvoted 2 times

Kp002 Most Recent 3 weeks, 5 days ago

Selected Answer: A

Publish data to Amazon Kinesis Data Streams. Use Kinesis Data Analytics to query the data.

- Kinesis Data Streams is designed for real-time data ingestion and can handle high throughput (1 MB/s is well within its capabilities).
- Data durability: Kinesis Data Streams stores data for up to 7 days (default is 24 hours), which helps prevent data loss during EC2 instance reboots.
- Kinesis Data Analytics allows you to run SQL queries on streaming data in near real-time, which satisfies the data science team's requirement.
- Scalable: Kinesis is designed to automatically scale with data volume, making it highly suitable for fluctuating ingestion

upvoted 1 times

Dantecito 4 months ago

Selected Answer: A

Also A because "The company's data science team wants to query ingested data in near-real time" and kinesis data analytics is doing that. With option B we query the data once it is in the data lake.

upvoted 1 times

Dantecito 4 months ago

Selected Answer: A

Option a for me. I will assume that the EC2 instance is responsible for placing the data into the correct database. Therefore, we just need to wait for the instance to reboot. Since we can store the data in Kinesis Data Streams for up to 7 days, we are covered in that regard.

For option B, we are forced to use Amazon Redshift, and the problem statement does not mention anything related to that. I would choose option A because it does not affect the original requirements, and it helps us avoid in-flight data loss as requested.

upvoted 1 times

wwwxxch 5 months, 1 week ago

Selected Answer: B

near-real time --> Kinesis Data Firehose

And retention day of Kinesis Data Streams cannot be longer than 365 days

upvoted 1 times

EllenLiu 5 months, 3 weeks ago

Selected Answer: A

A: focus on performing complex data processing without in-flight data lost, not mention data persistence

B: focus on data persist for later analysis

upvoted 1 times

LeonSauveterre 6 months, 1 week ago

Selected Answer: A

Instance store & ElastiCache are all temporary storages, which cannot address data loss. That rules out C & D.

B: Kinesis Data Firehose is optimized for batch processing rather than real-time querying. It can indeed deliver data to S3 or Redshift, but there's a good chance the delay between ingestion and query availability cannot meet the "near-real-time" requirement.

upvoted 1 times

Lin878 1 year ago

Selected Answer: B

https://aws.amazon.com/pm/kinesis/?gclid=CjwKCAjwvIWzBhAlEiwAHHWgvRQuJmBubZDnO2GasDWwc2iBapfVD6GBelgj2JV6qkldm-K_CmMzmxCdCwQAvD_BwE&trk=ee1218b7-7c10-4762-97df-274836a44566&sc_channel=ps&ef_id=CjwKCAjwvIWzBhAlEiwAHHWgvRQuJmBubZDnO2GasDWwc2iBapfVD6GBelgj2JV6qkldm-K_CmMzmxCdCwQAvD_BwE:G:s&s_kwid=AL!4422!3!651510255264!p!!g!!kinesis%20stream!19836376690!149589222920
upvoted 2 times

ray320x 1 year, 4 months ago

Option A is actually correct. The question ask for minimal data loss and that query of data should be near real time, not the ingestion. Kinesis data analytics is near real time.

Recent changes to Redshift actually make B correct as well, but A is also correct.
upvoted 2 times

dkw2342 1 year, 3 months ago

Streaming ingestion provides low-latency, high-speed ingestion of stream data from Amazon Kinesis Data Streams and Amazon Managed Streaming for Apache Kafka into an Amazon Redshift provisioned or Amazon Redshift Serverless materialized view.[1]

Option B mentions Kinesis Data Firehose (now just Firehose), so this won't work.

Option A is the correct answer.

[1]<https://docs.aws.amazon.com/redshift/latest/dg/materialized-view-streaming-ingestion.html>
upvoted 2 times

farnamjam 1 year, 4 months ago

Selected Answer: A

Comparison to other options:

B. Kinesis Data Firehose with Redshift: While Redshift is scalable, it doesn't offer real-time querying capabilities. Data needs to be loaded into Redshift from Firehose, introducing latency.

C. EC2 instance store with Kinesis Data Firehose and S3: Storing data in an EC2 instance store is not persistent and data will be lost during reboots. EBS volumes are more appropriate for persistent storage, but the architecture becomes more complex.

D. EBS volume with ElastiCache and Redis: While ElastiCache offers fast in-memory storage, it's not designed for high-volume data ingestion like 1 MB/s. It might struggle with scalability and persistence.

upvoted 3 times

Firdous586 1 year, 4 months ago

I don't understand why people are giving wrong information
in the QUESTION its clearly mentioned near Real Time
Kinesis Data Streams is for Real time

Where are Kinesis Datafirehose is for Near real time there for answer is B only
upvoted 5 times

Marco_St 1 year, 6 months ago

Selected Answer: A

Read the question: near real-time querying of data.... it is more about real-time data query once the data is ingested, It does not mention how long time the data needs to be stored. A is better option. B introduces delay of data buffer before it can be queried in redshift

upvoted 1 times

practice_makes_perfect 1 year, 6 months ago

Selected Answer: B

A is not correct because Kinesis can only store data up to 1 year. The solution need to support querying ALL data instead of "recent" data.

upvoted 3 times

pentium75 1 year, 5 months ago

Says who? They want to "query ingested data in near-real time", it does not say anything about historical data.

upvoted 2 times

Ruffyt 1 year, 6 months ago

A: is the solution for the company's requirements. Publishing data to Amazon Kinesis Data Streams can support ingestion rates as high as 1 MB/s and provide real-time data processing. Kinesis Data Analytics can query the ingested data in real-time with low latency, and the solution can scale as needed to accommodate increases in ingestion rates or querying needs. This solution also ensures minimal data loss in the event of an EC2 instance reboot since Kinesis Data Streams has a persistent data store for up to 7 days by default.

upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: A

Publish data to Amazon Kinesis Data Streams, Use Kinesis Data Analytics to query the data

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

- Provide near-real-time data ingestion into Kinesis Data Streams with the ability to handle the 1 MB/s ingestion rate. Data would be stored redundantly across shards.
- Enable near-real-time querying of the data using Kinesis Data Analytics. SQL queries can be run directly against the Kinesis data stream.
- Minimize data loss since data is replicated across shards. If an EC2 instance is rebooted, the data stream is still accessible.
- Scale seamlessly to handle varying ingestion and query rates.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #321

Topic 1

What should a solutions architect do to ensure that all objects uploaded to an Amazon S3 bucket are encrypted?

- A. Update the bucket policy to deny if the PutObject does not have an s3:x-amz-acl header set.
- B. Update the bucket policy to deny if the PutObject does not have an s3:x-amz-acl header set to private.
- C. Update the bucket policy to deny if the PutObject does not have an aws:SecureTransport header set to true.
- D. Update the bucket policy to deny if the PutObject does not have an x-amz-server-side-encryption header set. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

bdp123 **Highly Voted** 1 year, 9 months ago

Selected Answer: D

<https://aws.amazon.com/blogs/security/how-to-prevent-uploads-of-unencrypted-objects-to-amazon-s3/#:~:text=Solution%20overview>

upvoted 15 times

Grace83 1 year, 8 months ago

Thank you!

upvoted 1 times

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: D

The x-amz-server-side-encryption header is used to specify the encryption method that should be used to encrypt objects uploaded to an Amazon S3 bucket. By updating the bucket policy to deny if the PutObject does not have this header set, the solutions architect can ensure that all objects uploaded to the bucket are encrypted.

upvoted 6 times

awsgeek75 **Most Recent** 11 months, 1 week ago

Selected Answer: D

Related reading because (as of Jan 2023) S3 buckets have encryption enabled by default.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingServerSideEncryption.html>

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/using-server-side-encryption.html>

"If you require your data uploads to be encrypted using only Amazon S3 managed keys, you can use the following bucket policy. For example, the following bucket policy denies permissions to upload an object unless the request includes the x-amz-server-side-encryption header to request server-side encryption."

upvoted 4 times

kruasan 1 year, 7 months ago

To encrypt an object at the time of upload, you need to add a header called x-amz-server-side-encryption to the request to tell S3 to encrypt the object using SSE-C, SSE-S3, or SSE-KMS. The following code example shows a Put request using SSE-S3. <https://aws.amazon.com/blogs/security/how-to-prevent-uploads-of-unencrypted-objects-to-amazon-s3/>

upvoted 4 times

kruasan 1 year, 7 months ago

The other options would not enforce encryption:

A) Requiring an s3:x-amz-acl header does not mandate encryption. This header controls access permissions.

B) Requiring an s3:x-amz-acl header set to private also does not enforce encryption. It only enforces private access permissions.

C) Requiring an aws:SecureTransport header ensures uploads use SSL but does not specify that objects must be encrypted. Encryption is not required when using SSL transport.

upvoted 5 times

kruasan 1 year, 7 months ago

Selected Answer: D

To encrypt an object at the time of upload, you need to add a header called x-amz-server-side-encryption to the request to tell S3 to encrypt the object using SSE-C, SSE-S3, or SSE-KMS. The following code example shows a Put request using SSE-S3. <https://aws.amazon.com/blogs/security/how-to-prevent-uploads-of-unencrypted-objects-to-amazon-s3/>

upvoted 2 times

Sbbh 1 year, 8 months ago

Confusing question. It doesn't state clearly if the object needs to be encrypted at-rest or in-transit

upvoted 6 times

Guru4Cloud 1 year, 3 months ago

That's true

upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: D

I vote d

upvoted 2 times

LuckyAro 1 year, 9 months ago

Selected Answer: D

To ensure that all objects uploaded to an Amazon S3 bucket are encrypted, the solutions architect should update the bucket policy to deny any PutObject requests that do not have an x-amz-server-side-encryption header set. This will prevent any objects from being uploaded to the bucket unless they are encrypted using server-side encryption.

upvoted 4 times

jennyka76 1 year, 9 months ago

answer - D

upvoted 2 times

zTopic 1 year, 9 months ago

Selected Answer: D

Answer is D

upvoted 2 times

Neorem 1 year, 9 months ago

Selected Answer: D

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/amazon-s3-policy-keys.html>
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #322

Topic 1

A solutions architect is designing a multi-tier application for a company. The application's users upload images from a mobile device. The application generates a thumbnail of each image and returns a message to the user to confirm that the image was uploaded successfully.

The thumbnail generation can take up to 60 seconds, but the company wants to provide a faster response time to its users to notify them that the original image was received. The solutions architect must design the application to asynchronously dispatch requests to the different application tiers.

What should the solutions architect do to meet these requirements?

A. Write a custom AWS Lambda function to generate the thumbnail and alert the user. Use the image upload process as an event source to invoke the Lambda function.

B. Create an AWS Step Functions workflow. Configure Step Functions to handle the orchestration between the application tiers and alert the user when thumbnail generation is complete.

C. Create an Amazon Simple Queue Service (Amazon SQS) message queue. As images are uploaded, place a message on the SQS queue for thumbnail generation. Alert the user through an application message that the image was received.

Most Voted

D. Create Amazon Simple Notification Service (Amazon SNS) notification topics and subscriptions. Use one subscription with the application to generate the thumbnail after the image upload is complete. Use a second subscription to message the user's mobile app by way of a push notification after thumbnail generation is complete.

Correct Answer: C*Community vote distribution*

C (94%)

A (6%)

Comments**Steve_4542636** **Highly Voted** 1 year, 9 months ago**Selected Answer: C**

I've noticed there are a lot of questions about decoupling services and SQS is almost always the answer.
upvoted 33 times

alain_maza 1 week, 2 days ago

same that cloudfront
upvoted 1 times

Neha999 **Highly Voted** 1 year, 9 months ago

D
SNS fan out
upvoted 13 times

Manimgh **Most Recent** 1 month, 3 weeks ago

Selected Answer: C

C is correct -> take up to 60 seconds = SQS
upvoted 1 times

LoXoL 10 months ago

They don't look like real answers from the official exam...
upvoted 1 times

awsgeek75 10 months, 3 weeks ago

Selected Answer: C

Each option is badly worded:

A: "generate the thumbnail and alert the user" doesn't sound sequential so could alert the user during, before or after the thumbnail generation whichever way you interpret it.

B: this is sequential and won't alert until the steps are complete

D: Could work without with the risk of notification loss so C is better but this is also ok

upvoted 2 times

awsgeek75 11 months, 1 week ago

Selected Answer: C

Safe answer is C but B is so badly worded that it can mean anything to confuse people. Step functions to use tiers. What if on of the step is to inform to the user and move on to next step. Anyway, I'll chose C for the exam as it is cleaner.

upvoted 2 times

wsdasdasdqwda 1 year, 1 month ago

... asynchronously dispatch ... => Amazon SQS
upvoted 7 times

TariqKipkemei 1 year, 1 month ago

Selected Answer: C

Asynchronous, Decoupling = Amazon Simple Queue Service

upvoted 5 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

SQS is a fully managed message queuing service that can be used to decouple different parts of an application.

upvoted 2 times

Zox42 1 year, 8 months ago

Selected Answer: C

Answers B and D alert the user when thumbnail generation is complete. Answer C alerts the user through an application message that the image was received.

upvoted 5 times

Sbbh 1 year, 8 months ago

B:
Use cases for Step Functions vary widely, from orchestrating serverless microservices, to building data-processing pipelines, to defining a security-incident response. As mentioned above, Step Functions may be used for synchronous and asynchronous business processes.

business processes.

upvoted 1 times

AlessandraSAA 1 year, 9 months ago

why not B?

upvoted 4 times

Wael216 1 year, 9 months ago

Selected Answer: C

Creating an Amazon Simple Queue Service (SQS) message queue and placing messages on the queue for thumbnail generation can help separate the image upload and thumbnail generation processes.

upvoted 2 times

vindahake 1 year, 9 months ago

C

The key here is "a faster response time to its users to notify them that the original image was received." i.e user needs to be notified when image was received and not after thumbnail was created.

upvoted 3 times

AlmeroSenior 1 year, 9 months ago

Selected Answer: C

A looks like the best way , but its essentially replacing the mentioned app , that's not the ask

upvoted 2 times

Mickey321 1 year, 9 months ago

Selected Answer: A

<https://docs.aws.amazon.com/lambda/latest/dg/with-s3-tutorial.html>

upvoted 2 times

bdp123 1 year, 9 months ago

Selected Answer: C

C is the only one that makes sense

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #323

Topic 1

A company's facility has badge readers at every entrance throughout the building. When badges are scanned, the readers send a message over HTTPS to indicate who attempted to access that particular entrance.

A solutions architect must design a system to process these messages from the sensors. The solution must be highly available, and the results must be made available for the company's security team to analyze.

Which system architecture should the solutions architect recommend?

A. Launch an Amazon EC2 instance to serve as the HTTPS endpoint and to process the messages. Configure the EC2 instance to save the results to an Amazon S3 bucket.

B. Create an HTTPS endpoint in Amazon API Gateway. Configure the API Gateway endpoint to invoke an AWS Lambda function to process the messages and save the results to an Amazon DynamoDB table. **Most Voted**

C. Use Amazon Route 53 to direct incoming sensor messages to an AWS Lambda function. Configure the Lambda function to process the messages and save the results to an Amazon DynamoDB table.

D. Create a gateway VPC endpoint for Amazon S3. Configure a Site-to-Site VPN connection from the facility network to the VPC so that sensor data can be written directly to an S3 bucket by way of the VPC endpoint.

Correct Answer: B

Community vote distribution

B (100%)

Comments

kruasan **Highly Voted** 1 year, 1 month ago

Selected Answer: B

- Option A would not provide high availability. A single EC2 instance is a single point of failure.
- Option B provides a scalable, highly available solution using serverless services. API Gateway and Lambda can scale automatically, and DynamoDB provides a durable data store.
- Option C would expose the Lambda function directly to the public Internet, which is not a recommended architecture. API Gateway provides an abstraction layer and additional features like access control.
- Option D requires configuring a VPN to AWS which adds complexity. It also saves the raw sensor data to S3, rather than processing it and storing the results.

upvoted 22 times

TariqKipkemei Highly Voted 8 months ago

Selected Answer: B

Highly available = Serverless

The readers send a message over HTTPS = HTTPS endpoint in Amazon API Gateway

Process these messages from the sensors = AWS Lambda function

upvoted 8 times

Guru4Cloud Most Recent 9 months ago

Selected Answer: B

The correct answer is B. Create an HTTPS endpoint in Amazon API Gateway. Configure the API Gateway endpoint to invoke an AWS Lambda function to process the messages and save the results to an Amazon DynamoDB table.

Here are the reasons why:

API Gateway is a highly scalable and available service that can be used to create and expose RESTful APIs.

Lambda is a serverless compute service that can be used to process events and data.

DynamoDB is a NoSQL database that can be used to store data in a scalable and highly available way.

upvoted 4 times

Steve_4542636 1 year, 3 months ago

Selected Answer: B

I vote B

upvoted 2 times

KZM 1 year, 3 months ago

It is option "B"

Option "B" can provide a system with highly scalable, fault-tolerant, and easy to manage.

upvoted 2 times

LuckyAro 1 year, 3 months ago

Selected Answer: B

Deploy Amazon API Gateway as an HTTPS endpoint and AWS Lambda to process and save the messages to an Amazon DynamoDB table. This option provides a highly available and scalable solution that can easily handle large amounts of data. It also integrates with other AWS services, making it easier to analyze and visualize the data for the security team.

upvoted 4 times

zTopic 1 year, 3 months ago

Selected Answer: B

B is Correct

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #324

Topic 1

A company wants to implement a disaster recovery plan for its primary on-premises file storage volume. The file storage volume is mounted from an Internet Small Computer Systems Interface (iSCSI) device on a local storage server. The file storage volume holds hundreds of terabytes (TB) of data.

The company wants to ensure that end users retain immediate access to all file types from the on-premises systems without experiencing latency.

Which solution will meet these requirements with the LEAST amount of change to the company's existing infrastructure?

- A. Provision an Amazon S3 File Gateway as a virtual machine (VM) that is hosted on premises. Set the local cache to 10 TB. Modify existing applications to access the files through the NFS protocol. To recover from a disaster, provision an Amazon EC2 instance and mount the S3 bucket that contains the files.
- B. Provision an AWS Storage Gateway tape gateway. Use a data backup solution to back up all existing data to a virtual tape library. Configure the data backup solution to run nightly after the initial backup is complete. To recover from a disaster, provision an Amazon EC2 instance and restore the data to an Amazon Elastic Block Store (Amazon EBS) volume from the volumes in the virtual tape library.
- C. Provision an AWS Storage Gateway Volume Gateway cached volume. Set the local cache to 10 TB. Mount the Volume Gateway cached volume to the existing file server by using iSCSI, and copy all files to the storage volume. Configure scheduled snapshots of the storage volume. To recover from a disaster, restore a snapshot to an Amazon Elastic Block Store (Amazon EBS) volume and attach the EBS volume to an Amazon EC2 instance.
- D. Provision an AWS Storage Gateway Volume Gateway stored volume with the same amount of disk space as the existing file storage volume. Mount the Volume Gateway stored volume to the existing file server by using iSCSI, and copy all files to the storage volume. Configure scheduled snapshots of the storage volume. To recover from a disaster, restore a snapshot to an Amazon Elastic Block Store (Amazon EBS) volume and attach the EBS volume to an Amazon EC2 instance.

Most Voted**Correct Answer: D***Community vote distribution*

D (83%)

C (17%)

Comments

Grace83 Highly Voted 2 years, 2 months ago

D is the correct answer

Volume Gateway CACHED Vs STORED

Cached = stores a subset of frequently accessed data locally

Stored = Retains the ENTIRE ("all file types") in on prem data centre

upvoted 28 times

netcj Highly Voted 1 year, 9 months ago

Selected Answer: D

"users retain immediate access to all file types"

immediate cannot be cached -> D

upvoted 6 times

dkw2342 Most Recent 1 year, 3 months ago

Bad question. No RTO/RPO, so impossible to properly answer. They probably want to hear option D.

Depending on RPO, option B is also an adequate solution (data remains immediately accessible without experiencing latency via existing infrastructure, backup to cloud for DR). Also, this option requires LESS changes to existing infra than A. Only argument against B is that VTLs are usually used for legacy DR solutions, not for new ones, where object storage such as S3 is usually supported natively.

upvoted 2 times

MrPCarrot 1 year, 3 months ago

Answer is C go argue somewhere.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

A,B are wrong types of gateways for hundreds of TB of data that needs immediate access on-prem. C limits to 10TB. D provides access to all the files.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: D

"Immediate access to all file types from the on-premises systems without experiencing latency" requirement is not met by C. Also the solution is meant for DR purposes, the primary storage for the data should remain on premises.

upvoted 3 times

daniel1 1 year, 7 months ago

Selected Answer: C

From chatGPT4

Considering the requirements of minimal infrastructure change, immediate file access, and low-latency, Option C: Provisioning an AWS Storage Gateway Volume Gateway (cached volume) with a 10 TB local cache, seems to be the most fitting solution. This setup aligns with the existing iSCSI setup and provides a local cache for low-latency access, while also configuring scheduled snapshots for disaster recovery. In the event of a disaster, restoring a snapshot to an Amazon EBS volume and attaching it to an Amazon EC2 instance as described in this option would align with the recovery objective.

upvoted 1 times

pentium75 1 year, 5 months ago

ChatGPT is wrong. "Immediate access to all file types from the on-premises systems without experiencing latency" needs "stored volume" type. With "cached volume" not all data will be available locally.

upvoted 9 times

LoXoL 1 year, 3 months ago

pentium75 is right.

upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: D

End users retain immediate access to all file types = Volume Gateway stored volume
upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

dddddddd
upvoted 3 times

alexandercamachop 2 years ago

Selected Answer: D

Correct answer is Volume Gateway Stored which keeps all data on premises.
To have immediate access to the data. Cached is for frequently accessed data only.
upvoted 3 times

omoakin 2 years ago

CCCCCCCCCCCCCCCC
upvoted 1 times

24b2e9e 11 months, 3 weeks ago

The stored volume configuration stores the entire data set on-premises and asynchronously backs up the data to AWS. The cached volume configuration stores recently accessed data on-premises, and the remaining data is stored in Amazon S3 -that is why D is right
upvoted 2 times

lucdt4 2 years ago

Selected Answer: D

D is the correct answer
Volume Gateway CACHED Vs STORED
Cached = stores a data recently at local
Stored = Retains the ENTIRE ("all file types") in on prem data centre
upvoted 2 times

rushi0611 2 years, 1 month ago

Selected Answer: D

In the cached mode, your primary data is written to S3, while retaining your frequently accessed data locally in a cache for low-latency access.
In the stored mode, your primary data is stored locally and your entire dataset is available for low-latency access while asynchronously backed up to AWS.
Reference: <https://aws.amazon.com/storagegateway/faqs/>
Good luck.
upvoted 3 times

kruasan 2 years, 1 month ago

Selected Answer: D

It is stated the company wants to keep the data locally and have DR plan in cloud. It points directly to the volume gateway
upvoted 2 times

UnluckyDucky 2 years, 2 months ago

Selected Answer: D

"The company wants to ensure that end users retain immediate access to all file types from the on-premises systems "

D is the correct answer.
upvoted 3 times

CapJackSparrow 2 years, 2 months ago

Selected Answer: C

all file types, NOT all files. Volume mode can not cache 100TBs.

upvoted 3 times

eddie5049 2 years, 1 month ago

<https://docs.aws.amazon.com/storagegateway/latest/vgw/StorageGatewayConcepts.html>

Stored volumes can range from 1 GiB to 16 TiB in size and must be rounded to the nearest GiB. Each gateway configured for stored volumes can support up to 32 volumes and a total volume storage of 512 TiB (0.5 PiB).

upvoted 2 times

MssP 2 years, 2 months ago

all file types. Cached only save the most frequently or lastest accessed. If you didn't access any type for a long time, you will not cache it -> No immediate access

upvoted 4 times

pentium75 1 year, 5 months ago

Also the solution is meant for DR purposes, it's not like they need more storage or so.

upvoted 2 times

WhericanIstart 2 years, 2 months ago

Selected Answer: D

"The company wants to ensure that end users retain immediate access to all file types from the on-premises systems "

This points to stored volumes..

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #325

Topic 1

A company is hosting a web application from an Amazon S3 bucket. The application uses Amazon Cognito as an identity provider to authenticate users and return a JSON Web Token (JWT) that provides access to protected resources that are stored in another S3 bucket.


Upon deployment of the application, users report errors and are unable to access the protected content. A solutions architect must resolve this issue by providing proper permissions so that users can access the protected content.

Which solution meets these requirements?

- A. Update the Amazon Cognito identity pool to assume the proper IAM role for access to the protected content. **Most Voted**
- B. Update the S3 ACL to allow the application to access the protected content.
- C. Redeploy the application to Amazon S3 to prevent eventually consistent reads in the S3 bucket from affecting the ability of users to access the protected content.
- D. Update the Amazon Cognito pool to use custom attribute mappings within the identity pool and grant users the proper permissions to access the protected content.

Correct Answer: A

Community vote distribution



Comments

alexandercamachop **Highly Voted** 1 year ago

Selected Answer: A

To resolve the issue and provide proper permissions for users to access the protected content, the recommended solution is:

A. Update the Amazon Cognito identity pool to assume the proper IAM role for access to the protected content.

Explanation:

Amazon Cognito provides authentication and user management services for web and mobile applications. In this scenario, the application is using Amazon Cognito as an identity provider to authenticate users and obtain JSON Web Tokens (JWTs). The JWTs are used to access protected resources stored in another S3 bucket.

To grant users access to the protected content, the proper IAM role needs to be assumed by the identity pool in Amazon Cognito.

By updating the Amazon Cognito identity pool with the appropriate IAM role, users will be authorized to access the protected content in the S3 bucket.

upvoted 16 times

alexandercamachop 1 year ago

Option B is incorrect because updating the S3 ACL (Access Control List) will only affect the permissions of the application, not the users accessing the content.

Option C is incorrect because redeploying the application to Amazon S3 will not resolve the issue related to user access permissions.

Option D is incorrect because updating custom attribute mappings in Amazon Cognito will not directly grant users the proper permissions to access the protected content.

upvoted 12 times

LuckyAro **Highly Voted** 1 year, 3 months ago

Selected Answer: A

A is the best solution as it directly addresses the issue of permissions and grants authenticated users the necessary IAM role to access the protected content.

A suggests updating the Amazon Cognito identity pool to assume the proper IAM role for access to the protected content. This is a valid solution, as it would grant authenticated users the necessary permissions to access the protected content.

upvoted 6 times

Marco_St **Most Recent** 6 months ago

Selected Answer: A

IAM role is assigned to IAM users or groups or assumed by AWS service. So IAM role is given to AWS Cognito service which provides temporary AWS credentials to authenticated users. so technically When a user is authenticated by Cognito, they receive temporary credentials based on the IAM role tied to the Cognito identity pool. If this IAM role has permissions to access certain S3 buckets or objects, the authenticated user will be able to access those resources as allowed by the role. This service is used under the hood by Cognito to provide these temporary credentials. The credentials are limited in time and scope based on the permissions defined in the IAM role.

upvoted 2 times

Guru4Cloud 9 months ago

Selected Answer: A

A. Update the Amazon Cognito identity pool to assume the proper IAM role for access to the protected content.

upvoted 3 times

Abrar2022 12 months ago

Selected Answer: A

Services access other services via IAM Roles. Hence why updating AWS Cognito identity pool to assume proper IAM Role is the right solution.

upvoted 2 times

shanwford 1 year, 2 months ago

Selected Answer: A

Amazon Cognito identity pools assign your authenticated users a set of temporary, limited-privilege credentials to access your AWS resources. The permissions for each user are controlled through IAM roles that you create.

<https://docs.aws.amazon.com/cognito/latest/developerguide/role-based-access-control.html>

upvoted 3 times

Brak 1 year, 3 months ago

Selected Answer: D

A makes no sense - Cognito is not accessing the S3 resource. It just returns the JWT token that will be attached to the S3 request.

D is the right answer, using custom attributes that are added to the JWT and used to grant permissions in S3. See <https://docs.aws.amazon.com/cognito/latest/developerguide/using-attributes-for-access-control-policy-example.html> for an example.

upvoted 2 times

Abhineet9148232 1 year, 3 months ago

But even D requires setting up the permissions as bucket policy (as show in the shared example) which includes higher overhead than managing permissions attached to specific roles.

upvoted 3 times

asoli 1 year, 2 months ago

A says "Identity Pool"

According to AWS: "With an identity pool, your users can obtain temporary AWS credentials to access AWS services, such as Amazon S3 and DynamoDB."

So, answer is A

upvoted 4 times

Steve_4542636 1 year, 3 months ago

Selected Answer: A

Services access other services via IAM Roles.

upvoted 2 times

jennyka76 1 year, 3 months ago

ANSWER - A

<https://docs.aws.amazon.com/cognito/latest/developerguide/tutorial-create-identity-pool.html>

You have to create an custom role such as read-only

upvoted 5 times

zTopic 1 year, 3 months ago

Selected Answer: A

Answer is A

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #326

Topic 1

An image hosting company uploads its large assets to Amazon S3 Standard buckets. The company uses multipart upload in parallel by using S3 APIs and overwrites if the same object is uploaded again. For the first 30 days after upload, the objects will be accessed frequently. The objects will be used less frequently after 30 days, but the access patterns for each object will be inconsistent. The company must optimize its S3 storage costs while maintaining high availability and resiliency of stored assets.

Which combination of actions should a solutions architect recommend to meet these requirements? (Choose two.)

- A. Move assets to S3 Intelligent-Tiering after 30 days. **Most Voted**
- B. Configure an S3 Lifecycle policy to clean up incomplete multipart uploads. **Most Voted**
- C. Configure an S3 Lifecycle policy to clean up expired object delete markers.
- D. Move assets to S3 Standard-Infrequent Access (S3 Standard-IA) after 30 days.
- E. Move assets to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 30 days.

Correct Answer: AB*Community vote distribution***Comments****Neha999** **Highly Voted** 2 years, 3 months ago

AB

A : Access Pattern for each object inconsistent, Infrequent Access

B : Deleting Incomplete Multipart Uploads to Lower Amazon S3 Costs

upvoted 26 times

TungPham **Highly Voted** 2 years, 3 months ago**Selected Answer: AB**

B because Abort Incomplete Multipart Uploads Using S3 Lifecycle => <https://aws.amazon.com/blogs/aws-cloud-financial-management/discovering-and-deleting-incomplete-multipart-uploads-to-lower-amazon-s3-costs/>

A because The objects will be used less frequently after 30 days, but the access patterns for each object will be inconsistent => random access => S3 Intelligent-Tiering

upvoted 16 times

LeonSauveterre Most Recent 6 months, 1 week ago

Selected Answer: AB

Inconsistent access patterns: A is correct, and D is out (could work but less ideal and costs more).

C: "object delete markers" only apply to versioned buckets.

upvoted 1 times

ChymKuBoy 10 months, 3 weeks ago

Selected Answer: AB

AB for sure

upvoted 3 times

bujuman 1 year, 2 months ago

Selected Answer: BD

If we consider these statements:

1. For the first 30 days after upload, the objects will be accessed frequently
 2. The objects will be used less frequently after 30 days, but the access patterns for each object will be inconsistent
 3. The company must optimize its S3 storage costs while maintaining high availability and resiliency of stored assets.
 4. The company uses multipart upload in parallel by using S3 APIs and overwrites if the same object is uploaded again.
- Statements 1 and 2 could be completed with option D and not A because data is infrequently accessed only after 30 days. Due to usage of multipart upload, to meet requirement regarding cost optimization, option B will be used to clean up buckets uncompleted file parts (statements 3 & 4).

upvoted 4 times

NayeraB 1 year, 3 months ago

Selected Answer: AD

Because A & D address the main ask, there's no mention of cost optimization.

upvoted 2 times

NayeraB 1 year, 3 months ago

Facepalm It does ask for reducing the cost, A&B it is!

upvoted 3 times

NayeraB 1 year, 3 months ago

Selected Answer: AC

Because A & C address the main ask, there's no mention of cost optimization.

upvoted 1 times

NayeraB 1 year, 3 months ago

Not C :D, I meant to say A&D. Added another vote for that one.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: AB

A as the access pattern for each object is inconsistent so let AWS AWS do the handling.

B deals with multi-part duplication issues and saves money by deleting incomplete uploads

C No mention of deleted object so this is a distractor

D The objects will be accessed in unpredictable pattern so can't use this

E Not HA compliant

upvoted 4 times

awsgeek75 1 year, 5 months ago

Also, don't be confused by 30 days. The question has tricky wording: "The objects will be used less frequently after 30 days, but the access patterns for each object will be inconsistent"

It does NOT say that objects will be accessed less frequently after 30 days. It says the access is unpredictable which means it could go up or down. Don't make assumptions.

upvoted 5 times

awsgeek75 1 year, 5 months ago

pentium75 1 year, 5 months ago

Selected Answer: AB

C is nonsense

E does not meet the "high availability and resiliency" requirement

B is obvious (incomplete multipart uploads consume space -> cost money)

The tricky part is A vs. D. However, 'inconsistent access patterns' are the primary use case for Intelligent-Tiering. There are probably objects that will never be accessed and that would be moved to Glacier Instant Retrieval by Intelligent-Tiering, thus the overall cost would be lower than with D.

upvoted 4 times

osmk 1 year, 5 months ago

bd <https://docs.aws.amazon.com/AmazonS3/latest/userguide/storage-class-intro.html#sc-infreq-data-access> => S3 Standard-IA objects

are resilient to the loss of an Availability Zone. This storage class offers greater availability and resiliency than the S3 One Zone-IA class

upvoted 2 times

raymondfekry 1 year, 5 months ago

Selected Answer: AB

I wouldn't go with D since "the access patterns for each object will be inconsistent.", so we cannot move all assets to IA

upvoted 2 times

Marco_St 1 year, 6 months ago

Selected Answer: AB

inconsistent access pattern brings more sense to use Intelligent-Tiering after 30 days which also covers infrequent access.

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: AB

A. Move assets to S3 Intelligent-Tiering after 30 days.

B. Configure an S3 Lifecycle policy to clean up incomplete multipart uploads.

upvoted 2 times

vini15 1 year, 10 months ago

should be A and B

upvoted 2 times

MrAWSAssociate 1 year, 11 months ago

Selected Answer: BD

Option A has not been mentioned for resiliency in S3, check the page:

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/disaster-recovery-resiliency.html>

Therefore, I am with B & D choices.

upvoted 2 times

pentium75 1 year, 5 months ago

Intelligent-Tiering just moves to Standard-IA or Glacier Instant Access based on access patterns. This does not affect resiliency.

upvoted 2 times

alexandercamachop 2 years ago

Selected Answer: AB

A. Move assets to S3 Intelligent-Tiering after 30 days.

B. Configure an S3 Lifecycle policy to clean up incomplete multipart uploads.

Explanation:

A. Moving assets to S3 Intelligent-Tiering after 30 days: This storage class automatically analyzes the access patterns of objects and moves them between frequent access and infrequent access tiers. Since the objects will be accessed frequently for the first 30 days, storing them in the frequent access tier during that period optimizes performance. After 30 days, when the access patterns become inconsistent, S3 Intelligent-Tiering will automatically move the objects to the infrequent access tier, reducing

storage costs.

B. Configuring an S3 Lifecycle policy to clean up incomplete multipart uploads: Multipart uploads are used for large objects, and incomplete multipart uploads can consume storage space if not cleaned up. By configuring an S3 Lifecycle policy to clean up incomplete multipart uploads, unnecessary storage costs can be avoided.

upvoted 2 times

antropaws 2 years ago

Selected Answer: AD

AD.

B makes no sense because multipart uploads overwrite objects that are already uploaded. The question never says this is a problem.

upvoted 2 times

VellaDevil 1 year, 11 months ago

Questions says to optimize cost and if incomplete multipart uploads are not aborted it will still use capacity on S3 Bucket thus increase unnecessary cost.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #327

Topic 1

A solutions architect must secure a VPC network that hosts Amazon EC2 instances. The EC2 instances contain highly sensitive data and run in a private subnet. According to company policy, the EC2 instances that run in the VPC can access only approved third-party software repositories on the internet for software product updates that use the third party's URL. Other internet traffic must be blocked.

Which solution meets these requirements?

A. Update the route table for the private subnet to route the outbound traffic to an AWS Network Firewall firewall.

Configure domain list rule groups. **Most Voted**


B. Set up an AWS WAF web ACL. Create a custom set of rules that filter traffic requests based on source and destination IP address range sets.

C. Implement strict inbound security group rules. Configure an outbound rule that allows traffic only to the authorized software repositories on the internet by specifying the URLs.

D. Configure an Application Load Balancer (ALB) in front of the EC2 instances. Direct all outbound traffic to the ALB. Use a URL-based rule listener in the ALB's target group for outbound access to the internet.

Correct Answer: A

Community vote distribution



Comments

Bhawesh **Highly Voted** 1 year, 9 months ago

Selected Answer: A

Correct Answer A. Send the outbound connection from EC2 to Network Firewall. In Network Firewall, create stateful outbound rules to allow certain domains for software patch download and deny all other domains.

<https://docs.aws.amazon.com/network-firewall/latest/developerguide/suricata-examples.html#suricata-example-domain-filtering>

upvoted 16 times

Guru4Cloud 1 year, 3 months ago

Option A uses a network firewall which is overkill for instance-level rules.

upvoted 2 times

UnluckyDucky Highly Voted 1 year, 9 months ago

Selected Answer: A

Can't use URLs in outbound rule of security groups. URL Filtering screams Firewall.

upvoted 12 times

TheFivePips Most Recent 9 months, 2 weeks ago

Selected Answer: A

Security Groups operate at the transport layer (Layer 4) of the OSI model and are primarily concerned with controlling traffic based on IP addresses, ports, and protocols. They do not have the capability to inspect or filter traffic based on URLs.

The solution to restrict outbound internet traffic based on specific URLs typically involves using a proxy or firewall that can inspect the application layer (Layer 7) of the OSI model, where URL information is available.

AWS Network Firewall operates at the network and application layers, allowing for more granular control, including the ability to inspect and filter traffic based on domain names or URLs.

By configuring domain list rule groups in AWS Network Firewall, you can specify which URLs are allowed for outbound traffic. This option is more aligned with the requirement of allowing access to approved third-party software repositories based on their URLs.

upvoted 6 times

awsgeek75 11 months, 1 week ago

Selected Answer: A

<https://aws.amazon.com/network-firewall/features/>

"Web filtering:

AWS Network Firewall supports inbound and outbound web filtering for unencrypted web traffic. For encrypted web traffic, Server Name Indication (SNI) is used for blocking access to specific sites. SNI is an extension to Transport Layer Security (TLS) that remains unencrypted in the traffic flow and indicates the destination hostname a client is attempting to access over HTTPS. In addition, **AWS Network Firewall can filter fully qualified domain names (FQDN).**"

Always use an AWS product if the advertisement meets the use case.

upvoted 4 times

farnamjam 11 months, 2 weeks ago

Selected Answer: A

AWS Network Firewall

- Protect your entire Amazon VPC
- From Layer 3 to Layer 7 protection
- Any direction, you can inspect

Traffic filtering: Allow, drop, or alert for the traffic that matches the rules, • Active flow inspection to intrusion prevention

upvoted 2 times

Subhrangsu 11 months, 3 weeks ago

D not possible?

upvoted 1 times

awsgeek75 11 months, 1 week ago

ALB is for inbound traffic. D is not possible as it is suggesting to direct OUTBOUND traffic.

upvoted 4 times

Cyberkayu 11 months, 4 weeks ago

Selected Answer: A

AWS network firewall is stateful, providing control and visibility to Layer 3-7 network traffic, thus cover the application too

upvoted 2 times

TariqKipkemei 1 year, 1 month ago

Selected Answer: A

Just tried on the console to set up an outbound rule, and URLs cannot be used as a destination. I will opt for A.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

Implement strict inbound security group rules
Configure an outbound security group rule to allow traffic only to the approved software repository URLs
The key points:

Highly sensitive EC2 instances in private subnet that can access only approved URLs
Other internet access must be blocked
Security groups act as a firewall at the instance level and can control both inbound and outbound traffic.
upvoted 2 times

pentium75 11 months, 2 weeks ago

Security Groups work with CIDR ranges, not URLs.
upvoted 4 times

kelvintoys93 1 year, 5 months ago

Isn't private subnet not connectible to internet at all, unless with a NAT gateway?
upvoted 4 times

VeseljkoD 1 year, 9 months ago

Selected Answer: A

We can't specify URL in outbound rule of security group. Create free tier AWS account and test it.
upvoted 2 times

Leo301 1 year, 9 months ago

Selected Answer: C

CCCCCCCCCCC
upvoted 1 times

pentium75 11 months, 2 weeks ago

Security Groups with IP ranges, not URLs
upvoted 2 times

Brak 1 year, 9 months ago

It can't be C. You cannot use URLs in the outbound rules of a security group.
upvoted 3 times

johnmcclane78 1 year, 9 months ago

Option C is the best solution to meet the requirements of this scenario. Implementing strict inbound security group rules that only allow traffic from approved sources can help secure the VPC network that hosts Amazon EC2 instances. Additionally, configuring an outbound rule that allows traffic only to the authorized software repositories on the internet by specifying the URLs will ensure that only approved third-party software repositories can be accessed from the EC2 instances. This solution does not require any additional AWS services and can be implemented using VPC security groups.

Option A is not the best solution as it involves the use of AWS Network Firewall, which may introduce additional operational overhead. While domain list rule groups can be used to block all internet traffic except for the approved third-party software repositories, this solution is more complex than necessary for this scenario.
upvoted 2 times

pentium75 11 months, 2 weeks ago

How do you use a Security Group to allow access to https://server.com/repoa while denying access to https://server.com/repoa ? Security Groups work with IP ranges.
upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: C

In the security group, only allow inbound traffic originating from the VPC. Then only allow outbound traffic with a whitelisted IP address. The question asks about blocking EC2 instances, which is best for security groups since those are at the EC2 instance level. A network firewall is at the VPC level, which is not what the question is asking to protect.
upvoted 1 times

pentium75 11 months, 2 weeks ago

Security Groups work with IP ranges, not URLs.
upvoted 1 times

Theodorz 1 year, 9 months ago

Is Security Group able to allow a specific URL? According to https://docs.aws.amazon.com/vpc/latest/userguide/VPC_SecurityGroups.html, I cannot find such description.
upvoted 2 times

KZM 1 year, 9 months ago

I am confused that It seems both options A and C are valid solutions.
upvoted 3 times

Mia2009687 1 year, 5 months ago

I think C is in private subnet. Even with security group, it could not go public to download the software.
upvoted 1 times

Zohx 1 year, 9 months ago

Same here - why is C not a valid option?
upvoted 2 times

pentium75 11 months, 2 weeks ago

Because you want to filter based on URLs, not IP ranges.
upvoted 2 times

Karlos99 1 year, 9 months ago

And it is easier to do it at the VPC level
upvoted 2 times

Karlos99 1 year, 9 months ago

And it is easier to do it at the level
upvoted 2 times

ruqui 1 year, 6 months ago

C is not valid. Security groups can allow only traffic from specific ports and/or IPs, you can't use an URL. Correct answer is A
upvoted 3 times

jennyka76 1 year, 9 months ago

Answer - A
<https://aws.amazon.com/premiumsupport/knowledge-center/ec2-al1-al2-update-yum-without-internet/>
upvoted 6 times

asoli 1 year, 8 months ago

Although the answer is A, the link you provided here is not related to this question.
The information about "Network Firewall" and how it can help this issue is here:
<https://docs.aws.amazon.com/network-firewall/latest/developerguide/suricata-examples.html#suricata-example-domain-filtering>

(thanks to "@Bhawesh" to provide the link in their answer)
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #328

Topic 1

A company is hosting a three-tier ecommerce application in the AWS Cloud. The company hosts the website on Amazon S3 and integrates the website with an API that handles sales requests. The company hosts the API on three Amazon EC2 instances behind an Application Load Balancer (ALB). The API consists of static and dynamic front-end content along with backend workers that process sales requests asynchronously.

The company is expecting a significant and sudden increase in the number of sales requests during events for the launch of new products.

What should a solutions architect recommend to ensure that all the requests are processed successfully?

- A. Add an Amazon CloudFront distribution for the dynamic content. Increase the number of EC2 instances to handle the increase in traffic.
- B. Add an Amazon CloudFront distribution for the static content. Place the EC2 instances in an Auto Scaling group to launch new instances based on network traffic.
- C. Add an Amazon CloudFront distribution for the dynamic content. Add an Amazon ElastiCache instance in front of the ALB to reduce traffic for the API to handle.
- D. Add an Amazon CloudFront distribution for the static content. Add an Amazon Simple Queue Service (Amazon SQS) queue to receive requests from the website for later processing by the EC2 instances. **Most Voted**

Correct Answer: D*Community vote distribution*

D (73%)

B (27%)

Comments**Steve_4542636** **Highly Voted** 2 years, 3 months ago**Selected Answer: B**

The auto-scaling would increase the rate at which sales requests are "processed", whereas a SQS will ensure messages don't get lost. If you were at a fast food restaurant with a long line with 3 cash registers, would you want more cash registers or longer ropes to handle longer lines? Same concept here.

upvoted 22 times

AdamVinas 3 months, 3 weeks ago

AdumTiguo 6 months, 6 weeks ago

they want to process it successfully not faster (SQS) is the way...

upvoted 1 times

Chef_couincouin 1 year, 7 months ago

ensure that all the requests are processed successfully? doesn't mean more quickly

upvoted 4 times

lizzard812 2 years, 2 months ago

Hell true: I'd rather combine the both options: a SQS + auto-scaled bound to the length of the queue.

upvoted 10 times

joechen2023 1 year, 11 months ago

As an architecture, it is not possible to add more backend workers (it is part of the HR and boss's job, not for architecture design the solution). So when the demand surge, the only correct choice is to buffer them using SQS so that workers can take their time to process it successfully

upvoted 2 times

Abhineet9148232 **Highly Voted** 2 years, 1 month ago

Selected Answer: D

B doesn't fit because Auto Scaling alone does not guarantee that all requests will be processed successfully, which the question clearly asks for.

D ensures that all messages are processed.

upvoted 17 times

ChhatwaniB **Most Recent** 1 month, 4 weeks ago

Selected Answer: D

There is no scaling based on network traffic. since the backend worker works asynchronously SQS should be right here

upvoted 1 times

samadal 8 months ago

The problem states that the application consists of "static and dynamic front-end content." Static content typically includes cacheable resources such as HTML, CSS, and image files. Therefore, from this statement, one can infer that caching static content using CloudFront would improve performance. In other words, the mention of "static content" in the problem itself leads to the conclusion that CloudFront should be added for static content.

Additionally, the problem mentions "asynchronously processed backend workers." Asynchronous processing is well-suited for services like SQS, which can improve efficiency by handling dynamic requests that do not require immediate processing. The mention of "successfully processing all requests" also suggests that SQS is needed to ensure that all requests are handled properly.

Therefore, the correct answer is D.

upvoted 3 times

Adinas_ 1 year, 3 months ago

Selected Answer: B

Important question to answer D. Can you connect the website with SQS directly? How do you control access to who can put messages to SQS? I have never seen such a situation it has to be at least behind API gateway. So that conclusion brings me to answer B, application also can process async everything without SQS.

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: D

I chose D because I love SQS! These questions are hammering SQS in every solution as a "protagonist" that saves the day. AC are clearly useless

B can work but D is better because of SQS being better than EC2 scaling. The other part is that backend workers process the request asynchronously therefore a queue is better.

upvoted 4 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

A and C don't solve anything so ignore them.

Between B and D, D guarantees the scaling via SQS and order processing. B can also do that but it is not guaranteed that EC2 scaling will work to process the order.

As usual, I suspect that this "brain dump" may be missing critical wording to differentiate between the options so read carefully in the exam.

upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: D

There are two components that we need

* Frontend: Hosted on S3, performance can be increased with CloudFront

* Backend: There's no reason to process all the orders instantly, so we should decouple the processing from the API which we do with SQS

Thus D, CloudFront + SQS

upvoted 8 times

pentium75 1 year, 5 months ago

And as others said, B might speed up the processing or reduce the number of lost orders, but we need to make sure that "ALL requests are processed successfully", NOT that "less requests are lost".

upvoted 3 times

Marco_St 1 year, 6 months ago

Selected Answer: D

I picked B before I read D option. Read the question again, it concerns asynchronous processing of sales requests, Option D seems to align more closely with the requirements. So the requirement is ensuring all requests are processed successfully which means no request would be missed. So D is better option

upvoted 4 times

wsdadasdqwdaw 1 year, 7 months ago

Amazon SQS will make sure that the requests are stored and didn't get lost. After that the workers asynchronously will process the requests. I would go for D

upvoted 4 times

TariqKipkemei 1 year, 8 months ago

Technically both option B and D would work. But, there's a need to process requests asynchronously, hence decoupling, hence Amazon SQS. I will settle with option D.

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

D is correct.

upvoted 3 times

antropaws 2 years ago

Selected Answer: D

D is correct.

upvoted 2 times

kruasan 2 years, 1 month ago

Selected Answer: D

An SQS queue acts as a buffer between the frontend (website) and backend (API). Web requests can dump messages into the queue at a high throughput, then the queue handles delivering those messages to the API at a controlled rate that it can sustain. This prevents the API from being overwhelmed.

upvoted 3 times

kruasan 2 years, 1 month ago

Options A and B would help by scaling out more instances, however, this may not scale quickly enough and still risks

Options A and B would help by scaling out more instances, however, this may not scale quickly enough and still risks overwhelming the API. Caching parts of the dynamic content (option C) may help but does not provide the buffering mechanism that a queue does.

upvoted 1 times

seifshendy99 2 years, 1 month ago

Selected Answer: D

D make sens

upvoted 2 times

kraken21 2 years, 2 months ago

Selected Answer: D

D makes more sense

upvoted 2 times

kraken21 2 years, 2 months ago

There is no clarity on what the asynchronous process is but D makes more sense if we want to process all requests successfully. The way the question is worded it looks like the msgs->SQS>ELB/Ec2. This ensures that the messages are processed but may be delayed as the load increases.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #329

Topic 1

A security audit reveals that Amazon EC2 instances are not being patched regularly. A solutions architect needs to provide a solution that will run regular security scans across a large fleet of EC2 instances. The solution should also patch the EC2 instances on a regular schedule and provide a report of each instance's patch status.

Which solution will meet these requirements?

- A. Set up Amazon Macie to scan the EC2 instances for software vulnerabilities. Set up a cron job on each EC2 instance to patch the instance on a regular schedule.
- B. Turn on Amazon GuardDuty in the account. Configure GuardDuty to scan the EC2 instances for software vulnerabilities. Set up AWS Systems Manager Session Manager to patch the EC2 instances on a regular schedule.
- C. Set up Amazon Detective to scan the EC2 instances for software vulnerabilities. Set up an Amazon EventBridge scheduled rule to patch the EC2 instances on a regular schedule.
- D. Turn on Amazon Inspector in the account. Configure Amazon Inspector to scan the EC2 instances for software vulnerabilities. Set up AWS Systems Manager Patch Manager to patch the EC2 instances on a regular schedule. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

elearningtakai **Highly Voted** 1 year, 8 months ago

Selected Answer: D

Amazon Inspector is a security assessment service that automatically assesses applications for vulnerabilities or deviations from best practices. It can be used to scan the EC2 instances for software vulnerabilities. AWS Systems Manager Patch Manager can be used to patch the EC2 instances on a regular schedule. Together, these services can provide a solution that meets the requirements of running regular security scans and patching EC2 instances on a regular schedule. Additionally, Patch Manager can provide a report of each instance's patch status.

upvoted 10 times

LuckyAro **Highly Voted** 1 year, 9 months ago

Selected Answer: D

Amazon Inspector is a security assessment service that helps improve the security and compliance of applications deployed on

Amazon Web Services (AWS). It automatically assesses applications for vulnerabilities or deviations from best practices. Amazon Inspector can be used to identify security issues and recommend fixes for them. It is an ideal solution for running regular security scans across a large fleet of EC2 instances.

AWS Systems Manager Patch Manager is a service that helps you automate the process of patching Windows and Linux instances. It provides a simple, automated way to patch your instances with the latest security patches and updates. Patch Manager helps you maintain compliance with security policies and regulations by providing detailed reports on the patch status of your instances.

upvoted 5 times

awsgeek75 Most Recent 11 months, 1 week ago

Selected Answer: D

A handy reference page for such questions is:
<https://aws.amazon.com/products/security/>
Amazon Inspector = vulnerability detection = patching
<https://aws.amazon.com/inspector/>

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

dddddddddd

upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: D

Inspector is for EC2 instances and network accessibility of those instances
<https://portal.tutorialsdojo.com/forums/discussion/difference-between-security-hub-detective-and-inspector/>
upvoted 2 times

TungPham 1 year, 9 months ago

Selected Answer: D

Amazon Inspector for EC2
https://aws.amazon.com/vi/inspector/faqs/?nc1=f_ls
Amazon system manager Patch manager for automates the process of patching managed nodes with both security-related updates and other types of updates.

<http://webcache.googleusercontent.com/search?q=cache:FbFTc6XKycwJ:https://medium.com/aws-architech/use-case-aws-inspector-vs-guardduty-3662bf80767a&hl=vi&gl=kr&strip=1&vwsr=0>
upvoted 3 times

jennyka76 1 year, 9 months ago

answer - D
<https://aws.amazon.com/inspector/faqs/>
upvoted 3 times

Neha999 1 year, 9 months ago

D as AWS Systems Manager Patch Manager can patch the EC2 instances.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #330

Topic 1


A company is planning to store data on Amazon RDS DB instances. The company must encrypt the data at rest.

What should a solutions architect do to meet this requirement?

- A. Create a key in AWS Key Management Service (AWS KMS). Enable encryption for the DB instances. **Most Voted**
- B. Create an encryption key. Store the key in AWS Secrets Manager. Use the key to encrypt the DB instances.
- C. Generate a certificate in AWS Certificate Manager (ACM). Enable SSL/TLS on the DB instances by using the certificate.
- D. Generate a certificate in AWS Identity and Access Management (IAM). Enable SSL/TLS on the DB instances by using the certificate.

Correct Answer: A

Community vote distribution



Comments

awsgeek75 11 months, 1 week ago

Selected Answer: A

A: Enable encryption

B: KMS is for storage and doesn't directly integrate to DB without further work

C and D are for data encryption in transit not at rest

upvoted 4 times

awsgeek75 10 months, 3 weeks ago

Actually, D is total nonsense and no idea what it is saying

upvoted 1 times

robpalacios1 1 year ago

Selected Answer: A

KMS only generates and manages encryption keys. That's it. That's all it does. It's a fundamental service that you as well as other AWS Services (like Secrets Manager) use it to encrypt or decrypt.

Key Management Service. Secrets Manager is for database connection strings.

upvoted 3 times

invoted 4 times

upvoted 1 times

antropaws 1 year, 6 months ago

OK, but why not B???

upvoted 1 times

aaroncelestin 1 year, 3 months ago

KMS only generates and manages encryption keys. That's it. That's all it does. It's a fundamental service that you as well as other AWS Services (like Secrets Manager) use it to encrypt or decrypt.

Secrets Manager stores actual secrets like passwords, pass phrases, and anything else you want encrypted. SM uses KMS to encrypt its secrets, it would be circular to get an encryption key from KMS to use SM to encrypt the encryption key.

upvoted 5 times

SkyZeroZx 1 year, 7 months ago

Selected Answer: A

ANSWER - A

upvoted 1 times

datz 1 year, 8 months ago

Selected Answer: A

A for sure

upvoted 2 times

PRASAD180 1 year, 9 months ago

A is 100% Crt

upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: A

Key Management Service. Secrets Manager is for database connection strings.

upvoted 4 times

LuckyAro 1 year, 9 months ago

Selected Answer: A

A is the correct solution to meet the requirement of encrypting the data at rest.

To encrypt data at rest in Amazon RDS, you can use the encryption feature of Amazon RDS, which uses AWS Key Management Service (AWS KMS). With this feature, Amazon RDS encrypts each database instance with a unique key. This key is stored securely by AWS KMS. You can manage your own keys or use the default AWS-managed keys. When you enable encryption for a DB instance, Amazon RDS encrypts the underlying storage, including the automated backups, read replicas, and snapshots.

upvoted 4 times

bdp123 1 year, 9 months ago

Selected Answer: A

AWS Key Management Service (KMS) is used to manage the keys used to encrypt and decrypt the data.

upvoted 2 times

pbpally 1 year, 9 months ago

Selected Answer: A

Option A

upvoted 2 times

NolaHOla 1 year, 9 months ago

A. Create a key in AWS Key Management Service (AWS KMS). Enable encryption for the DB instances is the correct answer to encrypt the data at rest in Amazon RDS DB instances.

Amazon RDS provides multiple options for encrypting data at rest. AWS Key Management Service (KMS) is used to manage the keys used to encrypt and decrypt the data. Therefore, a solution architect should create a key in AWS KMS and enable

encryption for the DB instances to encrypt the data at rest.

upvoted 2 times

jennyka76 1 year, 9 months ago

ANSWER - A

<https://docs.aws.amazon.com/whitepapers/latest/efs-encrypted-file-systems/managing-keys.html>

upvoted 2 times

Bhawesh 1 year, 9 months ago

Selected Answer: A

A. Create a key in AWS Key Management Service (AWS KMS). Enable encryption for the DB instances.

<https://www.examttopics.com/discussions/amazon/view/80753-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 3 times

upvoted 1 times

wsdasdasdqwda 1 year, 1 month ago

$(15/8) = 1.875$ MB/s
 $1.875 \text{ MB/s} \times 0.7 = 1.3125$ (70% NW utilization) MB/s
 $1.3125 \text{ MB/s} \times 3600 = 4725$ MB (MB per 1 hour)
 $4725 \times 24 = 113400$ MB per 1 full day (24h)
 $113400 \times 30 = 3402000$ MB for 30 days
 $3402000 / 1024 = 3322.265625$ GB for 30 days
 $3322.265625 / 1024 \sim 3.24$ TB for 30 days => not enough for NW => Snowball which is A
upvoted 4 times

TariqKipkemei 1 year, 1 month ago

Selected Answer: A

I wont try to think to much about it, AWS Snowball was designed for this
upvoted 4 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: A

° 15 Mbps bandwidth with 70% max utilization limits the effective bandwidth to 10.5 Mbps or 1.31 MB/s.
° 20 TB of data at 1.31 MB/s would take approximately 193 days to transfer over the network. ° This far exceeds the 30 day requirement.
° AWS Snowball provides a physical storage device that can be shipped to the data center. Up to 80 TB can be loaded onto a Snowball device and shipped back to AWS.
This allows the 20 TB of data to be transferred much faster by shipping rather than over the limited network bandwidth.
° Snowball uses tamper-resistant enclosures and 256-bit encryption to keep the data secure during transit.
° The data can be imported into Amazon S3 or Amazon Glacier once the Snowball is received by AWS.
upvoted 5 times

UnluckyDucky 1 year, 8 months ago

Selected Answer: B

$10 \text{ MB/s} \times 86,400 \text{ seconds per day} \times 30 \text{ days} = 25,920,000 \text{ MB}$ or approximately 25.2 TB

That's how much you can transfer with a 10 Mbps link (roughly 70% of the 15 Mbps connection).

With a consistent connection of 8~ Mbps, and 30 days, you can upload 20 TB of data.

My math says B, my brain wants to go with A. Take your pick.

upvoted 3 times

Zox42 1 year, 8 months ago

$15 \text{ Mbps} \times 0.7 = 1.3125 \text{ MB/s}$ and $1.3125 \times 86,400 \times 30 = 3,402,000 \text{ MB}$
Answer A is correct.
upvoted 3 times

Zox42 1 year, 8 months ago

3,402,000
upvoted 2 times

hozy_ 1 year, 4 months ago

How can 15×0.7 be 1.3125 LMAO
upvoted 1 times

hozy_ 1 year, 4 months ago

OMG it was Mbps! Not MBps. You are right! awesome!!!
upvoted 3 times

Bilalazure 1 year, 9 months ago

Selected Answer: A

Aws snowball

upvoted 3 times

PRASAD180 1 year, 9 months ago

A is 100% Crt

upvoted 2 times

LuckyAro 1 year, 9 months ago

Selected Answer: A

AWS Snowball

upvoted 2 times

pbpally 1 year, 9 months ago

Selected Answer: A

Option a

upvoted 2 times

jennyka76 1 year, 9 months ago

ANSWER - A

<https://docs.aws.amazon.com/snowball/latest/ug/whatissnowball.html>

upvoted 2 times

AWSSHA1 1 year, 9 months ago

Selected Answer: A

option A

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #332

Topic 1

A company needs to provide its employees with secure access to confidential and sensitive files. The company wants to ensure that the files can be accessed only by authorized users. The files must be downloaded securely to the employees' devices.


The files are stored in an on-premises Windows file server. However, due to an increase in remote usage, the file server is running out of capacity.

Which solution will meet these requirements?

- A. Migrate the file server to an Amazon EC2 instance in a public subnet. Configure the security group to limit inbound traffic to the employees' IP addresses.
- B. Migrate the files to an Amazon FSx for Windows File Server file system. Integrate the Amazon FSx file system with the on-premises Active Directory. Configure AWS Client VPN. **Most Voted**
- C. Migrate the files to Amazon S3, and create a private VPC endpoint. Create a signed URL to allow download.
- D. Migrate the files to Amazon S3, and create a public VPC endpoint. Allow employees to sign on with AWS IAM Identity Center (AWS Single Sign-On).

Correct Answer: B

Community vote distribution



Comments

elearningtakai **Highly Voted** 2 years, 2 months ago

Selected Answer: B

This solution addresses the need for secure access to confidential and sensitive files, as well as the increase in remote usage. Migrating the files to Amazon FSx for Windows File Server provides a scalable, fully managed file storage solution in the AWS Cloud that is accessible from on-premises and cloud environments. Integration with the on-premises Active Directory allows for a consistent user experience and centralized access control. AWS Client VPN provides a secure and managed VPN solution that can be used by employees to access the files securely.

upvoted 9 times

pentium75 **Highly Voted** 1 year, 5 months ago

Selected Answer: B