

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #364

Topic 1

A hospital is designing a new application that gathers symptoms from patients. The hospital has decided to use Amazon Simple Queue Service (Amazon SQS) and Amazon Simple Notification Service (Amazon SNS) in the architecture.


A solutions architect is reviewing the infrastructure design. Data must be encrypted at rest and in transit. Only authorized personnel of the hospital should be able to access the data.

Which combination of steps should the solutions architect take to meet these requirements? (Choose two.)

- A. Turn on server-side encryption on the SQS components. Update the default key policy to restrict key usage to a set of authorized principals.
- B. Turn on server-side encryption on the SNS components by using an AWS Key Management Service (AWS KMS) customer managed key. Apply a key policy to restrict key usage to a set of authorized principals. **Most Voted**
- C. Turn on encryption on the SNS components. Update the default key policy to restrict key usage to a set of authorized principals. Set a condition in the topic policy to allow only encrypted connections over TLS.
- D. Turn on server-side encryption on the SQS components by using an AWS Key Management Service (AWS KMS) customer managed key. Apply a key policy to restrict key usage to a set of authorized principals. Set a condition in the queue policy to allow only encrypted connections over TLS. **Most Voted**
- E. Turn on server-side encryption on the SQS components by using an AWS Key Management Service (AWS KMS) customer managed key. Apply an IAM policy to restrict key usage to a set of authorized principals. Set a condition in the queue policy to allow only encrypted connections over TLS.

Correct Answer: BD

Community vote distribution



Comments

awsgeek75 **Highly Voted** 10 months, 3 weeks ago

My god! Every other question is about SQS! I thought this was AWS Solution Architect test not "How to solve any problem in AWS using SQS" test!

upvoted 16 times

fkie4 Highly Voted 1 year, 8 months ago

Selected Answer: BD

read this:

<https://docs.aws.amazon.com/sns/latest/dg/sns-server-side-encryption.html>

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-server-side-encryption.html>

upvoted 15 times

Gooniegoogoo 1 year, 5 months ago

good call.. that confirms on that page:

Important

All requests to topics with SSE enabled must use HTTPS and Signature Version 4.

For information about compatibility of other services with encrypted topics, see your service documentation.

Amazon SNS only supports symmetric encryption KMS keys. You cannot use any other type of KMS key to encrypt your service resources. For help determining whether a KMS key is a symmetric encryption key, see Identifying asymmetric KMS keys.

upvoted 4 times

pentium75 Most Recent 11 months, 1 week ago

Selected Answer: BD

A and C involve 'updating the default key policy' which is not something you. Either you create a key policy, OR AWS assigns THE "default key policy".

E 'applies an IAM policy to restrict key usage to a set of authorized principals' which is not how IAM policies work. You can 'apply an IAM policy to restrict key usage', but it would be restricted to the principals who have the policy attached; you can't specify them in the policy.

Leaves B and D. That B lacks the TLS statement is irrelevant because "all requests to topics with SSE enabled must use HTTPS" anyway.

upvoted 6 times

dkw2342 9 months ago

Yes, BD is correct.

"All requests to queues with SSE enabled must use HTTPS and Signature Version 4." -> valid for SNS and SQS alike:

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-server-side-encryption.html>

<https://docs.aws.amazon.com/sns/latest/dg/sns-server-side-encryption.html>

"Set a condition in the queue policy to allow only encrypted connections over TLS." refers to the "aws:SecureTransport" condition, but it's actually redundant.

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: CD

Its only options C and D that covers encryption on transit, encryption at rest and a restriction policy.

upvoted 3 times

Lalo 1 year, 6 months ago

Answer is BD

SNS: AWS KMS, key policy, SQS: AWS KMS, Key policy

upvoted 4 times

luisgu 1 year, 6 months ago

Selected Answer: BD

"IAM policies you can't specify the principal in an identity-based policy because it applies to the user or role to which it is attached"

reference: https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/security_iam_service-with-iam.html

that excludes E

unvoted 2 times

upvoted 2 times

imvb88 1 year, 7 months ago

Selected Answer: CD

Encryption at transit = use SSL/TLS -> rule out A,B

Encryption at rest = encryption on components -> keep C, D, E

KMS always need a key policy, IAM is optional -> E out

-> C, D left, one for SNS, one for SQS. TLS: checked, encryption on components: checked

upvoted 4 times

Lalo 1 year, 6 months ago

Answer is BD

SNS: AWS KMS, key policy, SQS: AWS KMS, Key policy

upvoted 4 times

imvb88 1 year, 7 months ago

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-data-encryption.html>

You can protect data in transit using Secure Sockets Layer (SSL) or client-side encryption. You can protect data at rest by requesting Amazon SQS to encrypt your messages before saving them to disk in its data centers and then decrypt them when the messages are received.

<https://docs.aws.amazon.com/kms/latest/developerguide/key-policies.html>

A key policy is a resource policy for an AWS KMS key. Key policies are the primary way to control access to KMS keys. Every KMS key must have exactly one key policy. The statements in the key policy determine who has permission to use the KMS key and how they can use it. You can also use IAM policies and grants to control access to the KMS key, but every KMS key must have a key policy.

upvoted 2 times

MarkGerwich 1 year, 8 months ago

CD

B does not include encryption in transit.

upvoted 3 times

Bofi 1 year, 8 months ago

That was my objection toward option B. CD cover both encryption at Rest and Server-Side_Encryption

upvoted 1 times

MssP 1 year, 8 months ago

in transit is included in D. With C, not include encryption at rest.... Server-side will include it.

upvoted 1 times

Maximus007 1 year, 8 months ago

ChatGPT returned AD as a correct answer)

upvoted 1 times

cegama543 1 year, 8 months ago

Selected Answer: BE

B: To encrypt data at rest, we can use a customer-managed key stored in AWS KMS to encrypt the SNS components.

E: To restrict access to the data and allow only authorized personnel to access the data, we can apply an IAM policy to restrict key usage to a set of authorized principals. We can also set a condition in the queue policy to allow only encrypted connections over TLS to encrypt data in transit.

upvoted 3 times

Karlos99 1 year, 8 months ago

Selected Answer: BD

For a customer managed KMS key, you must configure the key policy to add permissions for each queue producer and consumer. <https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-key-management.html>

upvoted 4 times

[Removed] 1 year, 9 months ago

Selected Answer: BE

bebebe

upvoted 1 times

[Removed] 1 year, 9 months ago

bdbdbdbd

All KMS keys must have a key policy. IAM policies are optional.

upvoted 7 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #365

Topic 1

A company runs a web application that is backed by Amazon RDS. A new database administrator caused data loss by accidentally editing information in a database table. To help recover from this type of incident, the company wants the ability to restore the database to its state from 5 minutes before any change within the last 30 days.

Which feature should the solutions architect include in the design to meet this requirement?

- A. Read replicas
- B. Manual snapshots
- C. Automated backups **Most Voted**
- D. Multi-AZ deployments

Correct Answer: C

Community vote distribution



Comments

Uzbekistan 9 months, 1 week ago

Selected Answer: C

Amazon RDS provides automated backups, which can be configured to take regular snapshots of the database instance. By enabling automated backups and setting the retention period to 30 days, the company can ensure that it retains backups for up to 30 days. Additionally, Amazon RDS allows for point-in-time recovery within the retention period, enabling the restoration of the database to its state from any point within the last 30 days, including 5 minutes before any change. This feature provides the required capability to recover from accidental data loss incidents.

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

Automated backups allow you to recover your database to any point in time within your specified retention period, which can be up to 35 days. The recovery process creates a new Amazon RDS instance with a new endpoint, and the process takes time proportional to the size of the database. Automated backups are enabled by default and occur daily during the backup window. This feature provides an easy and convenient way to recover from data loss incidents such as the one described in the scenario.

upvoted 4 times

elarningtakai 1 year, 8 months ago

Selected Answer: C

Option C, Automated backups, will meet the requirement. Amazon RDS allows you to automatically create backups of your DB instance. Automated backups enable point-in-time recovery (PITR) for your DB instance down to a specific second within the retention period, which can be up to 35 days. By setting the retention period to 30 days, the company can restore the database to its state from up to 5 minutes before any change within the last 30 days.

upvoted 4 times

joechen2023 1 year, 5 months ago

I selected C as well, but still don't know how the automatic backup could have a copy 5 minutes before any change. AWS doc states "Automated backups occur daily during the preferred backup window. "

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_WorkingWithAutomatedBackups.html.

I think the answer maybe A, as read replica will be kept sync and then restore from the read replica. could an expert help?

upvoted 2 times

awsgeek75 11 months, 1 week ago

"the company wants the ability to restore the database to its state from 5 minutes before any change"

The automatic backup takes a backup every 5 minutes. This means it can restore the database to 5 minutes in the past.

upvoted 2 times

TheFivePips 9 months, 2 weeks ago

Automated backups enable point-in-time recovery (PITR) for your DB instance down to a specific second within the retention period, which can be up to 35 days

upvoted 2 times

gold4otas 1 year, 8 months ago

Selected Answer: C

C: Automated Backups

<https://aws.amazon.com/rds/features/backup/>

upvoted 3 times

WheracanIstart 1 year, 8 months ago

Selected Answer: C

Automated Backups...

upvoted 3 times

[Removed] 1 year, 9 months ago

Selected Answer: C

cccccccc

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #366

Topic 1

A company's web application consists of an Amazon API Gateway API in front of an AWS Lambda function and an Amazon DynamoDB database. The Lambda function handles the business logic, and the DynamoDB table hosts the data. The application uses Amazon Cognito user pools to identify the individual users of the application. A solutions architect needs to update the application so that only users who have a subscription can access premium content.

Which solution will meet this requirement with the LEAST operational overhead?

- A. Enable API caching and throttling on the API Gateway API.
- B. Set up AWS WAF on the API Gateway API. Create a rule to filter users who have a subscription.
- C. Apply fine-grained IAM permissions to the premium content in the DynamoDB table.
- D. Implement API usage plans and API keys to limit the access of users who do not have a subscription. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****Guru4Cloud** **Highly Voted** 1 year, 3 months ago**Selected Answer: D**

Implementing API usage plans and API keys is a straightforward way to restrict access to specific users or groups based on subscriptions. It allows you to control access at the API level and doesn't require extensive changes to your existing architecture. This solution provides a clear and manageable way to enforce access restrictions without complicating other parts of the application

upvoted 12 times

Rcosmos **Most Recent** 4 months, 2 weeks ago**Selected Answer: C**

Segurança robusta:

A lógica de controle de acesso é aplicada no nível do banco de dados, garantindo que usuários não autorizados não consigam acessar os dados premium, mesmo que tentem burlar o aplicativo.

upvoted 1 times

Ilzhekistan 9 months, 1 week ago

awsgeek75 11 months, 1 week ago

Selected Answer: C

Chat GPT said:

Option C, "Apply fine-grained IAM permissions to the premium content in the DynamoDB table," would likely involve the least operational overhead.

Here's why:

Granular Control: IAM permissions allow you to control access at a very granular level, including specific actions (e.g., GetItem, PutItem) on individual resources (e.g., DynamoDB tables).

Integration with Cognito: IAM policies can be configured to allow access based on the identity of the user authenticated through Cognito. You can create IAM roles or policies that grant access to users with specific attributes or conditions, such as having a subscription.

Minimal Configuration Changes: This solution primarily involves configuring IAM policies for access control in DynamoDB, which can be done with minimal changes to the existing application architecture.

upvoted 1 times

awsgeek75 11 months, 1 week ago

Selected Answer: C

C is correct as per the link and doc:

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-usage-plans.html#apigateway-usage-plans-best-practices>

D: API keys cannot be used to limit access and this can only be done via methods defined in above link

upvoted 2 times

awsgeek75 11 months, 1 week ago

Also, option A is for performance and not for security

option B, WAF cannot control access based on subscription without massive custom coding which will be a big operational overhead

upvoted 2 times

awsgeek75 10 months, 3 weeks ago

I had to chose D but must have clicked C incorrectly. It is D as my explanation is about D not C! C is the wrong answer.

upvoted 2 times

lipi0035 1 year ago

In the same document <https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-usage-plans.html> if you scroll down, it says `Don't use API keys for authentication or authorization to control access to your APIs. If you have multiple APIs in a usage plan, a user with a valid API key for one API in that usage plan can access all APIs in that usage plan. Instead, to control access to your API, use an IAM role, a Lambda authorizer, or an Amazon Cognito user pool.`

In the same document at the bottom, it says "If you're using a developer portal to publish your APIs, note that all APIs in a given usage plan are subscribable, even if you haven't made them visible to your customers."

I go with C

upvoted 2 times

awsgeek75 11 months, 1 week ago

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-usage-plans.html#apigateway-usage-plans-best-practices>

Correct link

upvoted 2 times

TariqKipkemei 1 year, 1 month ago

Selected Answer: D

After you create, test, and deploy your APIs, you can use API Gateway usage plans to make them available as product offerings for your customers. You can configure usage plans and API keys to allow customers to access selected APIs, and begin throttling requests to those APIs based on defined limits and quotas. These can be set at the API, or API method level.

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-usage-plans.html#:~:text=Creating%20and%20using-,usage%20plans,-with%20API%20keys>

upvoted 2 times

marufexplorer 1 year, 5 months ago

D

Option D involves implementing API usage plans and API keys. By associating specific API keys with users who have a valid subscription, you can control access to the premium content.

upvoted 2 times

kruasan 1 year, 7 months ago

Selected Answer: D

A. This would not actually limit access based on subscriptions. It helps optimize and control API usage, but does not address the core requirement.

B. This could work by checking user subscription status in the WAF rule, but would require ongoing management of WAF and increases operational overhead.

C. This is a good approach, using IAM permissions to control DynamoDB access at a granular level based on subscriptions. However, it would require managing IAM permissions which adds some operational overhead.

D. This option uses API Gateway mechanisms to limit API access based on subscription status. It would require the least amount of ongoing management and changes, minimizing operational overhead. API keys could be easily revoked/changed as subscription status changes.

upvoted 4 times

imvb88 1 year, 7 months ago

CD both possible but D is more suitable since it mentioned in

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-usage-plans.html>

A,B not relevant.

upvoted 3 times

elarningtakai 1 year, 8 months ago

Selected Answer: D

The solution that will meet the requirement with the least operational overhead is to implement API Gateway usage plans and API keys to limit access to premium content for users who do not have a subscription.

Option A is incorrect because API caching and throttling are not designed for authentication or authorization purposes, and it does not provide access control.

Option B is incorrect because although AWS WAF is a useful tool to protect web applications from common web exploits, it is not designed for authorization purposes, and it might require additional configuration, which increases the operational overhead.

Option C is incorrect because although IAM permissions can restrict access to data stored in a DynamoDB table, it does not provide a mechanism for limiting access to specific content based on the user subscription. Moreover, it might require a significant amount of additional IAM permissions configuration, which increases the operational overhead.

upvoted 3 times

klayytech 1 year, 8 months ago

Selected Answer: D

To meet the requirement with the least operational overhead, you can implement API usage plans and API keys to limit the access of users who do not have a subscription. This way, you can control access to your API Gateway APIs by requiring clients to submit valid API keys with requests. You can associate usage plans with API keys to configure throttling and quota limits on individual client accounts.

upvoted 3 times

techhb 1 year, 8 months ago

answer is D ,if looking for least overhead

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-usage-plans.html>

C will achieve it but operational overhead is high.

upvoted 3 times

quentin17 1 year, 8 months ago

Selected Answer: D

Both C&D are valid solution

According to ChatGPT:

"Applying fine-grained IAM permissions to the premium content in the DynamoDB table is a valid approach, but it requires more effort in managing IAM policies and roles for each user, making it more complex and adding operational overhead."

upvoted 2 times

Karlos99 1 year, 8 months ago

Selected Answer: D

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-usage-plans.html>

upvoted 4 times

[Removed] 1 year, 9 months ago

Selected Answer: C

cccccccc

upvoted 1 times

pentium75 11 months, 1 week ago

"Fine-grained permissions" for only two groups of users, hell no.

"IAM permissions" for customers, also no.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #367

Topic 1


A company is using Amazon Route 53 latency-based routing to route requests to its UDP-based application for users around the world. The application is hosted on redundant servers in the company's on-premises data centers in the United States, Asia, and Europe. The company's compliance requirements state that the application must be hosted on premises. The company wants to improve the performance and availability of the application.

What should a solutions architect do to meet these requirements?

- A. Configure three Network Load Balancers (NLBs) in the three AWS Regions to address the on-premises endpoints. Create an accelerator by using AWS Global Accelerator, and register the NLBs as its endpoints. Provide access to the application by using a CNAME that points to the accelerator DNS. **Most Voted**
- B. Configure three Application Load Balancers (ALBs) in the three AWS Regions to address the on-premises endpoints. Create an accelerator by using AWS Global Accelerator, and register the ALBs as its endpoints. Provide access to the application by using a CNAME that points to the accelerator DNS.
- C. Configure three Network Load Balancers (NLBs) in the three AWS Regions to address the on-premises endpoints. In Route 53, create a latency-based record that points to the three NLBs, and use it as an origin for an Amazon CloudFront distribution. Provide access to the application by using a CNAME that points to the CloudFront DNS.
- D. Configure three Application Load Balancers (ALBs) in the three AWS Regions to address the on-premises endpoints. In Route 53, create a latency-based record that points to the three ALBs, and use it as an origin for an Amazon CloudFront distribution. Provide access to the application by using a CNAME that points to the CloudFront DNS.

Correct Answer: A

Community vote distribution



Comments

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: A

NLBs allow UDP traffic (ALBs don't support UDP)
 Global Accelerator uses Anycast IP addresses and its global network to intelligently route users to the optimal endpoint
 Using NLBs as Global Accelerator endpoints provides improved availability and DDoS protection.
 upvoted 12 times

lucdt4 **Highly Voted** 2 years ago

Selected Answer: A

C - D: Cloudfront don't support UDP/TCP
B: Global accelerator don't support ALB
A is correct

upvoted 6 times

JA2018 **Most Recent** 6 months, 2 weeks ago

Selected Answer: C

CloudFront's origin can be on-premises sources.

Check this out:

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/DownloadDistS3AndCustomOrigins.html#concept_CustomOrigin

"A custom origin is an HTTP server, for example, a web server. The HTTP server can be an Amazon EC2 instance or an HTTP server that you host somewhere else. "

upvoted 1 times

sandordini 1 year, 1 month ago

Selected Answer: A

Non-HTTP, Massive performance: NLB, UDP: AWS Global Accelerator

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: A

Neither ALB (B+D) nor CloudFront (C+D) do support UDP.

upvoted 5 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: A

UDP = NLB and Global Accelerator

upvoted 4 times

live_reply_developers 1 year, 11 months ago

Selected Answer: A

NLB + GA support UDP/TCP

upvoted 4 times

Gooniegoogoo 1 year, 11 months ago

good reference <https://blog.cloudcraft.co/alb-vs-nlb-which-aws-load-balancer-fits-your-needs/>

upvoted 3 times

SkyZeroZx 2 years, 1 month ago

Selected Answer: A

UDP = NBL

UDP = GLOBAL ACCELERATOR

UPD NOT WORKING WITH CLOUDFRONT

ANS IS A

upvoted 4 times

MssP 2 years, 2 months ago

Selected Answer: A

More discussions at: <https://www.examttopics.com/discussions/amazon/view/51508-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 1 times

Grace83 2 years, 2 months ago

Why is C not correct - does anyone know?

upvoted 2 times

MssP 2 years, 2 months ago

It could be valid but I think A is better. Uses the AWS global network to optimize the path from users to applications, improving the performance of TCP and UDP traffic

upvoted 2 times

Shrestwt 2 years, 1 month ago

Latency based routing is already using in the application, so AWS global network will optimize the path from users to applications.

upvoted 2 times

FourOfAKind 2 years, 2 months ago

Selected Answer: A

UDP == NLB

Must be hosted on-premises != CloudFront

upvoted 4 times

imvb88 2 years, 1 month ago

actually CloudFront's origin can be on-premises. Source:

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/DownloadDistS3AndCustomOrigins.html#concept_CustomOrigin

"A custom origin is an HTTP server, for example, a web server. The HTTP server can be an Amazon EC2 instance or an HTTP server that you host somewhere else. "

upvoted 2 times

[Removed] 2 years, 3 months ago

Selected Answer: A

aaaaaaaa

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #368

Topic 1


A solutions architect wants all new users to have specific complexity requirements and mandatory rotation periods for IAM user passwords.

What should the solutions architect do to accomplish this?

- A. Set an overall password policy for the entire AWS account. **Most Voted**
- B. Set a password policy for each IAM user in the AWS account.
- C. Use third-party vendor software to set password requirements.
- D. Attach an Amazon CloudWatch rule to the Create_newuser event to set the password with the appropriate requirements.

Correct Answer: A

Community vote distribution



Comments

lostmagnet001 **Highly Voted** 1 year, 3 months ago

Selected Answer: A

i get confused, the question says "NEW" users... if you apply this password policy it would affect all the users in the AWS account....

upvoted 10 times

angel_marquina **Highly Voted** 1 year, 8 months ago

The question is for new users, answer A is not exact for that case.

upvoted 9 times

Rcosmos **Most Recent** 4 months, 2 weeks ago

Selected Answer: U

Por que as outras opções não são adequadas?

B. Definir uma política de senha para cada usuário do IAM:

O IAM não permite políticas de senha configuradas individualmente para usuários. As políticas de senha são aplicadas a nível da conta.

upvoted 1 times

LeonSauveterre 6 months, 1 week ago

Selected Answer: A

A: Overall password policy for the entire AWS account - affects new users immediately, but not old users immediately. Next time the old users change their passwords, the new rule will apply then.

B: Logically wrong. If you already have a user, then that user must've already have a password. So when you individually set a password policy for these users, they are only affected next time they change their passwords, just like I mentioned above about option A. Plus, this is really tiresome because this is user-wise, and there might be too many users, and the work you must do manually is increasing horribly.

upvoted 1 times

[Removed] 1 year ago

Selected Answer: B

Because its mentioned "all new users"

upvoted 1 times

[Removed] 1 year ago

Ignore the above, seems to be custom policy can help in the case, so A should be right

upvoted 1 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: A

You can set a custom password policy on your AWS account to specify complexity requirements and mandatory rotation periods for your IAM users' passwords. When you create or change a password policy, most of the password policy settings are enforced the next time your users change their passwords. However, some of the settings are enforced immediately.

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_credentials_passwords_account-policy.html#:~:text=Setting%20an%20account-,password%20policy,-for%20IAM%20users

upvoted 4 times

klayytech 2 years, 2 months ago

Selected Answer: A

To accomplish this, the solutions architect should set an overall password policy for the entire AWS account. This policy will apply to all IAM users in the account, including new users.

upvoted 4 times

WheracanIstart 2 years, 2 months ago

Selected Answer: A

Set overall password policy ...

upvoted 3 times

kampatra 2 years, 2 months ago

Selected Answer: A

A is correct

upvoted 2 times

[Removed] 2 years, 3 months ago

Selected Answer: A

aaaaaaa

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #369

Topic 1

A company has migrated an application to Amazon EC2 Linux instances. One of these EC2 instances runs several 1-hour tasks on a schedule. These tasks were written by different teams and have no common programming language. The company is concerned about performance and scalability while these tasks run on a single instance. A solutions architect needs to implement a solution to resolve these concerns.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use AWS Batch to run the tasks as jobs. Schedule the jobs by using Amazon EventBridge (Amazon CloudWatch Events).

Most Voted


B. Convert the EC2 instance to a container. Use AWS App Runner to create the container on demand to run the tasks as jobs.

C. Copy the tasks into AWS Lambda functions. Schedule the Lambda functions by using Amazon EventBridge (Amazon CloudWatch Events).

D. Create an Amazon Machine Image (AMI) of the EC2 instance that runs the tasks. Create an Auto Scaling group with the AMI to run multiple copies of the instance.

Correct Answer: A

Community vote distribution



Comments

fkie4 Highly Voted 1 year, 8 months ago

Selected Answer: C

question said "These tasks were written by different teams and have no common programming language", and key word "scalable". Only Lambda can fulfil these. Lambda can be done in different programming languages, and it is scalable
upvoted 10 times

wsdasdasdqwda 1 year, 1 month ago

AWS Batch - As a fully managed service, AWS Batch helps you to run batch computing workloads of any scale. AWS Batch automatically provisions compute resources and optimizes the workload distribution based on the quantity and scale of the workloads. With AWS Batch, there's no need to install or manage batch computing software, so you can focus your time on analyzing results and solving problems. <https://docs.aws.amazon.com/batch/latest/userguide/what-is-batch.html> ---> I am

emerging needs and solving problems like, scalability, security, safety, etc., what is essential for them voting for A, C would have been OK if the time was within 15 minutes.

upvoted 8 times

FourOfAKind 1 year, 8 months ago

But the question states "several 1-hour tasks on a schedule", and the maximum runtime for Lambda is 15 minutes, so it can't be A.

upvoted 32 times

FourOfAKind 1 year, 8 months ago

can't be C

upvoted 9 times

JTruong 11 months, 1 week ago

Lambda can only execute job under 15 mins* so C can't be the answer

upvoted 6 times

smgsi 1 year, 8 months ago

It's not because time limit of lambda is 15 minutes

upvoted 13 times

pentium75 Highly Voted 11 months, 1 week ago

"Running on a schedule" = Batch

Not C due Lambda < 15 min

Not D, auto-scaling doesn't make sense for things running on a schedule

upvoted 9 times

Manimgh Most Recent 1 month, 3 weeks ago

Selected Answer: A

no common programming language = Batch

upvoted 1 times

foha2012 10 months, 2 weeks ago

Selected Answer: D

Answer = D

"performance and scalability while these tasks run on a single instance" They gave me a legacy application and want it to autoscale for performance. They don't want it to run on a single EC2 instance. Shouldn't I make an AMI and provision multiple EC2 instances in an autoscaling group? I could put an ALB in front of it. I won't have to deal with "uncommon programming languages" inside the application... Just a thought..

upvoted 2 times

awsgeek75 11 months, 1 week ago

Selected Answer: A

AWS Batch is for jobs running at schedule on EC2. so option A

B is operational overhead

C Lambda is 15 mins max execution

D Scaling is not a requirement

upvoted 7 times

meowruki 1 year ago

Selected Answer: A

AWS Batch: AWS Batch is a fully managed service for running batch computing workloads. It dynamically provisions the optimal quantity and type of compute resources based on the volume and specific resource requirements of the batch jobs. It allows you to run tasks written in different programming languages with minimal operational overhead.

upvoted 4 times

hungta 1 year ago

Selected Answer: A

The task working for hour but lambda function timeout is 15 minutes. So vote A.

upvoted 2 times

upvoted 2 times

yodelin 1 year, 1 month ago

I know guys are stressed out trying to figure this exam out okay, but no matter what people say, with or without reasoning, at least put your mouth clean. Going like AAA is an issue, but talking shi* on him just because he didn't write down the reasoning is your fault.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

It can run heterogeneous workloads and tasks without needing to convert them to a common format. AWS Batch manages the underlying compute resources - no need to manage containers, Lambda functions or Auto Scaling groups.

upvoted 6 times

zjcorpuz 1 year, 4 months ago

AWS Lambda function can only be run for 15 mins

upvoted 2 times

jaydesai8 1 year, 5 months ago

Selected Answer: A

maximum runtime for Lambda is 15 minutes, hence A

upvoted 3 times

antropaws 1 year, 6 months ago

Selected Answer: A

I also go with A.

upvoted 2 times

omoakin 1 year, 6 months ago

C. Copy the tasks into AWS Lambda functions. Schedule the Lambda functions by using Amazon EventBridge (Amazon CloudWatch Events)

upvoted 1 times

ruqui 1 year, 6 months ago

wrong, Lambda maximum runtime is 15 minutes and the tasks run for an hour

upvoted 3 times

KMohsoe 1 year, 6 months ago

Selected Answer: A

B and D out!

A and C let's think!

AWS Lambda functions are time limited.

So, Option A

upvoted 2 times

lucdt4 1 year, 6 months ago

AAAAAAAAAAAAAAAAAAAA

because lambda only run within 15 minutes

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

Answer is A.

Could have been C but AWS Lambda functions can be only configured to run up to 15 minutes per execution. While the task in question need an 1hour to run,

upvoted 5 times

luisgu 1 year, 6 months ago

Selected Answer: D

question is asking for the LEAST operational overhead. With batch, you have to create the compute environment, create the job queue, create the job definition and create the jobs --> more operational overhead than creating an ASG
upvoted 2 times

pentium75 11 months, 1 week ago

Things 'running on a schedule' = Batch, not autoscaling
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #370

Topic 1

A company runs a public three-tier web application in a VPC. The application runs on Amazon EC2 instances across multiple Availability Zones. The EC2 instances that run in private subnets need to communicate with a license server over the internet. The company needs a managed solution that minimizes operational maintenance.

Which solution meets these requirements?

- A. Provision a NAT instance in a public subnet. Modify each private subnet's route table with a default route that points to the NAT instance.
- B. Provision a NAT instance in a private subnet. Modify each private subnet's route table with a default route that points to the NAT instance.
- C. Provision a NAT gateway in a public subnet. Modify each private subnet's route table with a default route that points to the NAT gateway. **Most Voted**
- D. Provision a NAT gateway in a private subnet. Modify each private subnet's route table with a default route that points to the NAT gateway.

Correct Answer: C

Community vote distribution

C (100%)

Comments

UnluckyDucky **Highly Voted** 1 year, 8 months ago

Selected Answer: C

"The company needs a managed solution that minimizes operational maintenance"

Watch out for NAT instances vs NAT Gateways.

As the company needs a managed solution that minimizes operational maintenance - NAT Gateway in a public subnet is the answer.

upvoted 9 times

lucdt4 **Highly Voted** 1 year, 6 months ago

C
NAT gateway can't deploy in a private subnet

NAT gateway can't deploy in a private subnet.
upvoted 5 times

von_himmlen Most Recent 7 months ago

C
<https://docs.aws.amazon.com/appstream2/latest/developerguide/managing-network-internet-NAT-gateway.html>
...and a NAT gateway in a public subnet.
upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

This meets the requirements for a managed, low maintenance solution for private subnets to access the internet:

NAT gateway provides automatic scaling, high availability, and fully managed service without admin overhead.
Placing the NAT gateway in a public subnet with proper routes allows private instances to use it for internet access.
Minimal operational maintenance compared to NAT instances.
upvoted 3 times

Guru4Cloud 1 year, 3 months ago

No good:
NAT instances (A, B) require more hands-on management.

Placing a NAT gateway in a private subnet (D) would not allow internet access.
upvoted 4 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

minimizes operational maintenance = NGW
upvoted 2 times

WhericanIstart 1 year, 8 months ago

Selected Answer: C

C..provision NGW in Public Subnet
upvoted 3 times

cegama543 1 year, 8 months ago

Selected Answer: C

ccccc is the best
upvoted 2 times

[Removed] 1 year, 9 months ago

Selected Answer: C

cccccccc
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #371

Topic 1

A company needs to create an Amazon Elastic Kubernetes Service (Amazon EKS) cluster to host a digital media streaming application. The EKS cluster will use a managed node group that is backed by Amazon Elastic Block Store (Amazon EBS) volumes for storage. The company must encrypt all data at rest by using a customer managed key that is stored in AWS Key Management Service (AWS KMS).

Which combination of actions will meet this requirement with the LEAST operational overhead? (Choose two.)

- A. Use a Kubernetes plugin that uses the customer managed key to perform data encryption.
- B. After creation of the EKS cluster, locate the EBS volumes. Enable encryption by using the customer managed key.
- C. Enable EBS encryption by default in the AWS Region where the EKS cluster will be created. Select the customer managed key as the default key. **Most Voted**
- D. Create the EKS cluster. Create an IAM role that has a policy that grants permission to the customer managed key. Associate the role with the EKS cluster. **Most Voted**
- E. Store the customer managed key as a Kubernetes secret in the EKS cluster. Use the customer managed key to encrypt the EBS volumes.

Correct Answer: CD*Community vote distribution***Comments****asoli** **Highly Voted** 2 years, 2 months ago**Selected Answer: CD**

<https://docs.aws.amazon.com/eks/latest/userguide/managed-node-groups.html#:~:text=encrypted%20Amazon%20EBS%20volumes%20without%20using%20a%20launch%20template%2C%20encrypt%20all%20new%20Amazon%20EBS%20volumes%20created%20in%20your%20account.>

upvoted 16 times

bujuman 1 year, 2 months ago

If you want to encrypt Amazon EBS volumes for your nodes, you can deploy the nodes using a launch template. To deploy managed nodes with encrypted Amazon EBS volumes without using a launch template, encrypt all new Amazon EBS volumes

created in your account. For more information, see Encryption by default in the Amazon EC2 User Guide for Linux Instances.

upvoted 3 times

imvb88 Highly Voted 2 years, 1 month ago

Selected Answer: BD

Quickly rule out A (which plugin? > overhead) and E because of bad practice

Among B,C,D: B and C are functionally similar > choice must be between B or C, D is fixed

Between B and C: C is out since it set default for all EBS volume in the region, which is more than required and even wrong, say what if other EBS volumes of other applications in the region have different requirement?

upvoted 12 times

NSA_Poker 1 year ago

(C) is correct; the EBS volumes of other applications in the region will not be affected bc an IAM role will limit the encryption key to the EKS cluster.

upvoted 2 times

srcntpc Most Recent 3 weeks, 4 days ago

Selected Answer: CD

Why C is Correct:

Enabling EBS encryption by default in the Region ensures that all newly created EBS volumes are encrypted automatically. You can specify a customer managed KMS key as the default encryption key.

Why D is Correct:

The EKS cluster (and its nodes) will need IAM permissions to use the customer managed KMS key for encryption/decryption. Associating the correct IAM role with the EKS cluster ensures this access.

upvoted 1 times

Yak_Yeti 1 month, 1 week ago

Selected Answer: CD

D is clear and as to B vs C? The trade-off is the potential for encrypting other volumes in the Region, but the manual effort and complexity of Option B outweigh this concern when considering operational efficiency.

upvoted 1 times

babayomi 7 months, 2 weeks ago

CD

This options are wrongly stated in my understanding. B is wrong because you can not encrypt an unencrypted created ebs volume. You need to take a snapshot of the volume, encrypt the snapshot ,creat new ebs volume from snapshot. You can also enable encryption during creation. So the statement is wrong. The next available option can only be C, except that in that account all ebs volumes created would be encrypted, this is also questionable. Because if another person create a new ebs volumes it's automatically encrypted.

upvoted 2 times

scaredSquirrel 9 months, 3 weeks ago

Selected Answer: CD

A and E are obvious nos. D is a shoo-in.

The difference between B&C is basically EBS encryption by default vs encryption. Encryption by default is by region, and encrypt everything in that region going forward, versus simple encryption is volume by volume, C is less operational overhead. Check doc & chatGPT.

upvoted 2 times

jjcode 1 year, 3 months ago

this one is going on my skip list

upvoted 7 times

Mahmoudddddddd 1 year, 2 months ago

Don't it came for me in my exam today xd

upvoted 4 times

JA2018 6 months, 2 weeks ago

Hi Mahmoudddddddd, can you share what were your chosen answers?

upvoted 1 times

jaswantn 1 year, 4 months ago

If question is giving a requirement related to a particular case and asking to encrypt all data at rest; it is clear that encryption is for this case only and not for other projects in entire region. so option B is more appropriate along with option D.

upvoted 2 times

frmrkc 1 year, 4 months ago

Selected Answer: CD

It says: 'The company must encrypt ALL data at rest', so there is nothing wrong with 'enabling EBS encryption by default' . C & D

upvoted 5 times

LeonSauveterre 6 months, 1 week ago

Exactly. Option B is out of the question. Not to mention option C barely has any operational overhead.

upvoted 1 times

upliftinghut 1 year, 5 months ago

Selected Answer: BD

B&D are correct. C is wrong because when you turn on encryption by default, AWS uses its own key while the requirement is using Customer key.

Detail is here: <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSEncryption.html#encryption-by-default>

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: BD

Not A (avoid 3rd party plugins when there are native services)

Not C ("encryption by default" would impact other services)

Not E (Keys belong in KMS, not in EKS cluster)

upvoted 4 times

awsgeek75 1 year, 5 months ago

"The company must encrypt all data at rest by using a customer managed key that is stored in AWS Key Management Service (AWS KMS)."

I am just a bit concerned that the question does not put any limits on not encrypting all the EBS by default in the account. Both B and C can work. C is a hack but it is definitely LEAST operational overhead. Also, we don't know if there are other services or not that may be impacted. What do you think?

upvoted 2 times

Marco_St 1 year, 5 months ago

Selected Answer: CD

EBS encryption is set regionally. AWS account is global but it does not mean EBS encryption is enable by default at account level. default EBS encryption is a regional setting within your AWS account. Enabling it in a specific region ensures that all new EBS volumes created in that region are encrypted by default, using either the default AWS managed key or a customer managed key that you specify.

upvoted 2 times

pentium75 1 year, 5 months ago

"Enabling it in a specific region ensures that all new EBS volumes created in that region are encrypted by default" which is not what we want. We want to encrypt the EBS volumes used by this EKS cluster, NOT "all new EBS volumes created in that region."

upvoted 2 times

maudsha 1 year, 7 months ago

Selected Answer: CD

IF you need to encrypt an unencrypted volume,

- Create an EBS snapshot of the volume
 - Encrypt the EBS snapshot (using copy)
 - Create new EBS volume from the snapshot (the volume will also be encrypted)
- so it has an operational overhead.

So assuming they won't use this account for anything else we can use C. Enable EBS encryption by default in the AWS Region where the EKS cluster will be created. Select the customer managed key as the default key.

upvoted 1 times

pentium75 1 year, 5 months ago

"Assuming they won't use this account for anything else" how could we assume that?

upvoted 2 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: CD

Option D is required wither way.

Technically both option B and C would work, but with B you would have to enable encryption node by node, while with option C provides a onetime action of enabling encryption on all nodes.

The requirement is the option with LEAST operational overhead.

upvoted 3 times

pentium75 1 year, 5 months ago

B created some deployment work, but NOT "operational (!) overhead" once it's deployed. C enables encryption by default for all new EBS volumes which is not what we want.

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: CD

These options allow EBS encryption with the customer managed KMS key with minimal operational overhead:

C) Setting the KMS key as the regional EBS encryption default automatically encrypts new EKS node EBS volumes.

D) The IAM role grants the EKS nodes access to use the key for encryption/decryption operations.

upvoted 1 times

jaydesai8 1 year, 11 months ago

Selected Answer: CD

C - enable EBS encryption by default in a region -<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSEncryption.html>

D - Provides key access permission just to the EKS cluster without changing broader IAM permissions

upvoted 1 times

pentium75 1 year, 5 months ago

We're not asked to enable EBS encryption by default.

upvoted 2 times

pedroso 1 year, 12 months ago

Selected Answer: BD

I was in doubt between B and C.

You can't "Enable EBS encryption by default in the AWS Region". Enable EBS encryption by default is only possible at Account level, not Region.

B is the right option once you can enable encryption on the EBS volume with KMS and custom KMS.

upvoted 2 times

antropaws 1 year, 11 months ago

Not accurate: "Encryption by default is a Region-specific setting":

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSEncryption.html#encryption-by-default>

upvoted 4 times

pentium75 1 year, 5 months ago

Still C is wrong because "encryption by default" is not what we want.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #372


Topic 1

A company wants to migrate an Oracle database to AWS. The database consists of a single table that contains millions of geographic information systems (GIS) images that are high resolution and are identified by a geographic code.

When a natural disaster occurs, tens of thousands of images get updated every few minutes. Each geographic code has a single image or row that is associated with it. The company wants a solution that is highly available and scalable during such events.

Which solution meets these requirements MOST cost-effectively?

- A. Store the images and geographic codes in a database table. Use Oracle running on an Amazon RDS Multi-AZ DB instance.
- B. Store the images in Amazon S3 buckets. Use Amazon DynamoDB with the geographic code as the key and the image S3 URL as the value. **Most Voted**
- C. Store the images and geographic codes in an Amazon DynamoDB table. Configure DynamoDB Accelerator (DAX) during times of high load.
- D. Store the images in Amazon S3 buckets. Store geographic codes and image S3 URLs in a database table. Use Oracle running on an Amazon RDS Multi-AZ DB instance.

Correct Answer: B*Community vote distribution***Comments****Wayne23Fang** **Highly Voted** 1 year, 9 months ago**Selected Answer: B**

Amazon prefers people to move from Oracle to its own services like DynamoDB and S3.
upvoted 16 times

Karlos99 **Highly Voted** 2 years, 2 months ago**Selected Answer: D**

The company wants a solution that is highly available and scalable

upvoted 9 times

EllenLiu 5 months, 3 weeks ago

another thing is the design of dynamoDB, "Use Amazon DynamoDB with the geographic code as the key and the image S3 URL as the value." geographic code can not be the only key for any image

upvoted 2 times

[Removed] 2 years, 2 months ago

But DynamoDB is also highly available and scalable

<https://aws.amazon.com/dynamodb/faqs/#:~:text=DynamoDB%20automatically%20scales%20throughput%20capacity,high%20availability%20and%20data%20durability.>

upvoted 5 times

pbpally 2 years, 1 month ago

Yes but has a size limit at 400kb so theoretically it could store images but it's not a plausible solution.

upvoted 2 times

ruqui 2 years ago

It doesn't matter the size limit of DynamoDB!!!! The images are saved in S3 buckets. Right answer is B

upvoted 9 times

jaydesai8 1 year, 11 months ago

but would it be easy and cost-effective to migrate Oracle (relational db) to (Dynamodb)NoSQL?

upvoted 5 times

pentium75 1 year, 5 months ago

Yes because it's a single table with two records, for which Oracle or any relation database has been a bad choice in the first place.

upvoted 6 times

upliftinghut Most Recent 1 year, 4 months ago

Selected Answer: B

DynamoDB with its HA and built-in scalability. The nature of the table also resonates with NoSQL than SQL DB such as Oracle. Only 1 table so migration is just a script from Oracle to DynamoDB

D is workable but more expensive with Oracle licenses and other setups for HA and scalability

upvoted 3 times

upliftinghut 1 year, 4 months ago

HA & built-in scalability of Amazon DynamoDB :

<https://aws.amazon.com/dynamodb/features/#:~:text=Amazon%20DynamoDB%20is%20a%20fully,for%20the%20most%20demanding%20applications.>

upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

A puts images in Oracle, not a good idea

C DAX is not going to help with images

D It is doable but RDS on multi AZ does not give you more performance or write scalability. It gives more availability and read scalability which is not required here.

B works as Geographic code is the key in DynamoDB and S3 image URL is the data so DynamoDB can handle tens of thousands such record and S3 can scale for writing

upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: B

They are currently using Oracle, but only for one simple table with a single key-value pair. This is a typical use case for a NoSQL database like DynamoDB (and whoever decided to use Oracle for this in the first place should be fired). Oracle is expensive as

database like DynamoDB (and whoever decided to use Oracle for this in the first place should be tired). Oracle is expensive as hell, so options A and D might work but are surely not cost-effective. C won't work because the images are too big for the database. Leaves B which would be the ideal solution and meet the availability and scalability requirements.

upvoted 7 times

wsdasdasdqwda 1 year, 7 months ago

For D - Oracle is not cheap as well. RDS with Oracle vs DynamoDB, I would go for pure AWS provided option. In each exam there is a lot of marketing => B

upvoted 3 times

jubolano 1 year, 7 months ago

Selected Answer: D

Cost effective, D

upvoted 2 times

awsgeek75 1 year, 4 months ago

How is Oracle more cost effective than other options?

upvoted 2 times

wsdasdasdqwda 1 year, 7 months ago

B or D, but the question is MOST cost-effectively DynamoDB is more expensive than RDS, I am going for D

upvoted 3 times

JA2018 6 months, 2 weeks ago

Key in sTEM states that the customer's current Oracle DB table is setup with a single key-value pair.

A typical use case for a NoSQL database like DynamoDB.

upvoted 1 times

gouranga45 1 year, 8 months ago

Selected Answer: B

Answer is B, DynamoDB is Highly available and scalable

upvoted 2 times

baba365 1 year, 8 months ago

A single table in a relational db can have items that are related ? e.g. 'select * from Faculty where department_id in (10, 20) and dept_name = AWS'.

In the sql query example above, * means all and Faculty is name of the table.

upvoted 2 times

Eminenza22 1 year, 9 months ago

Selected Answer: B

B option offers a cost-effective solution for storing and accessing high-resolution GIS images during natural disasters. Storing the images in Amazon S3 buckets provides scalable and durable storage, while using Amazon DynamoDB allows for quick and efficient retrieval of images based on geographic codes. This solution leverages the strengths of both S3 and DynamoDB to meet the requirements of high availability, scalability, and cost-effectiveness.

upvoted 2 times

cd93 1 year, 9 months ago

Selected Answer: B

What were the company thinking using the most expensive DB on the planet FOR ONE SINGLE TABLE???

Migrate a single table from SQL to NoSQL should be easy enough I guess...

upvoted 3 times

vini15 1 year, 10 months ago

Should be D.

the question says company wants to migrate oracle to AWS. Oracle is a relational db hence RDS makes more sense whereas Dynamodb is non relational db.

upvoted 3 times

pentium75 1 year, 5 months ago

But relational DB does not make sense for the use case. It's a single table.
upvoted 2 times

iBanan 1 year, 10 months ago

I hate these questions:) I can't choose between B and D
upvoted 6 times

ces_9999 1 year, 11 months ago

Guys the answer is B the oracle database only has one table without any relationships so why we should use a relational database in the first place, second we are storing the images in S3 not in the database why not use this alongside dynamo
upvoted 6 times

Kp88 1 year, 10 months ago

You can't do migration of Oracle to Dynmodb without SCT. I am not the DB guy but since its saying oracle I would go with D otherwise B makes more sense if a company is starting out from scratch.
upvoted 2 times

Kp88 1 year, 10 months ago

Actually now that I think about it , B sounds ok as well. Company just need to use SCT and that would be more cost effective.
upvoted 2 times

joehong 1 year, 11 months ago

Selected Answer: D

"A company wants to migrate an Oracle database to AWS"
upvoted 2 times

pentium75 1 year, 5 months ago

Yeah, per my understanding that doesn't implicate that the destination must be an Oracle database.
upvoted 2 times

secdgs 1 year, 11 months ago

D: Wronng
if you caluate License Oracle Database, It is not cost-effectively. Multi-AZ is not scalable and if you set scalable, you need more license for Oracle database.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #373

Topic 1

A company has an application that collects data from IoT sensors on automobiles. The data is streamed and stored in Amazon S3 through Amazon Kinesis Data Firehose. The data produces trillions of S3 objects each year. Each morning, the company uses the data from the previous 30 days to retrain a suite of machine learning (ML) models.


Four times each year, the company uses the data from the previous 12 months to perform analysis and train other ML models. The data must be available with minimal delay for up to 1 year. After 1 year, the data must be retained for archival purposes.

Which storage solution meets these requirements MOST cost-effectively?

- A. Use the S3 Intelligent-Tiering storage class. Create an S3 Lifecycle policy to transition objects to S3 Glacier Deep Archive after 1 year.
- B. Use the S3 Intelligent-Tiering storage class. Configure S3 Intelligent-Tiering to automatically move objects to S3 Glacier Deep Archive after 1 year.
- C. Use the S3 Standard-Infrequent Access (S3 Standard-IA) storage class. Create an S3 Lifecycle policy to transition objects to S3 Glacier Deep Archive after 1 year.
- D. Use the S3 Standard storage class. Create an S3 Lifecycle policy to transition objects to S3 Standard-Infrequent Access (S3 Standard-IA) after 30 days, and then to S3 Glacier Deep Archive after 1 year. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

UnluckyDucky **Highly Voted** 1 year, 8 months ago

Selected Answer: D

Access patterns is given, therefore D is the most logical answer.

Intelligent tiering is for random, unpredictable access.

upvoted 14 times

ealpuce 1 year, 6 months ago

You are missing: <<The data must be available with minimal delay for up to 1 year. After one year, the data must be retained for archival purposes.>> You are secure that data after 1 year is not accessible anymore.

upvoted 2 times

TariqKipkemei Highly Voted 1 year, 6 months ago

Selected Answer: D

First 30 days data accessed every morning = S3 Standard
Beyond 30 days data accessed quarterly = S3 Standard-Infrequent Access
Beyond 1 year data retained = S3 Glacier Deep Archive

upvoted 5 times

jjcode Most Recent 9 months, 4 weeks ago

I dont get how its A

1. Each morning, the company uses the data from the previous 30 days
2. Four times each year, the company uses the data from the previous 12 months to perform analysis and train other ML models
3. The data must be available with minimal delay for up to 1 year. After 1 year, the data must be retained for archival purposes

The data ingestion happens 4 times a year, that means that after the initial 30 days it still needs to be pulled 3 more times, why would you put the data in standard infrequent if you were going to use it 3 more times and speed is a requirement? Makes more sense to put it in S3 standard, or intelligent then straight to glacier.

upvoted 2 times

upliftinghut 10 months, 2 weeks ago

Selected Answer: D

Clear access pattern. data in Standard-Infrequent Access is for data requires rapid access when needed

upvoted 2 times

awsgeek75 11 months, 1 week ago

Selected Answer: D

A and B, Intelligent Tiering cannot be configured. It is managed by AWS.
C SIA does not allow immediate access for "each morning"
D is best for 30 day standard access, SIA after 30 days and archive after 1 year

upvoted 4 times

pentium75 11 months, 2 weeks ago

Selected Answer: D

See reasoning below, just accidentally voted A

upvoted 2 times

pentium75 11 months, 2 weeks ago

Selected Answer: A

The data is used every day (typical use case for Standard) for 30 days, for the remaining 12 months it is used 3 or 4 times (typical use case for IA), after 12 months it is not used at all but must be kept (typical use case for Glacier Deep Archive).

upvoted 2 times

pentium75 11 months, 2 weeks ago

Sorry, D!!!!!!!!! Not A!!!! D!

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

This option optimizes costs while meeting the data access requirements:

Store new data in S3 Standard for first 30 days of frequent access
Transition to S3 Standard-IA after 30 days for infrequent access up to 1 year
Archive to Glacier Deep Archive after 1 year for long-term archival

upvoted 3 times

ealpuche 1 year, 6 months ago

Selected Answer: A

Option A meets the requirements most cost-effectively. The S3 Intelligent-Tiering storage class provides automatic tiering of objects between the S3 Standard and S3 Standard-Infrequent Access (S3 Standard-IA) tiers based on changing access patterns, which helps optimize costs. The S3 Lifecycle policy can be used to transition objects to S3 Glacier Deep Archive after 1 year for archival purposes. This solution also meets the requirement for minimal delay in accessing data for up to 1 year. Option B is not cost-effective because it does not include the transition of data to S3 Glacier Deep Archive after 1 year. Option C is not the best solution because S3 Standard-IA is not designed for long-term archival purposes and incurs higher storage costs. Option D is also not the most cost-effective solution as it transitions objects to the S3 Standard-IA tier after 30 days, which is unnecessary for the requirement to retrain the suite of ML models each morning using data from the previous 30 days.

upvoted 1 times

pentium75 11 months, 2 weeks ago

I can't follow. The data is used every day (typical use case for Standard) for 30 days, for the remaining 12 months it is used 3 or 4 times (typical use case for IA), after 12 months it is not used at all but must be kept (typical use case for Glacier Deep Archive).

upvoted 3 times

KAUS2 1 year, 8 months ago

Selected Answer: D

Agree with UnluckyDucky , the correct option is D

upvoted 2 times

fkie4 1 year, 8 months ago

Selected Answer: D

Should be D. see this:

<https://www.examttopics.com/discussions/amazon/view/68947-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 3 times

Nithin1119 1 year, 9 months ago

Selected Answer: B

Bbbbbbbbbb

upvoted 1 times

fkie4 1 year, 8 months ago

hello!!??

upvoted 2 times

[Removed] 1 year, 9 months ago

Selected Answer: D

ddddddd

upvoted 4 times

[Removed] 1 year, 9 months ago

D because:

- First 30 days- data access every morning (predictable and frequently) – S3 standard
- After 30 days, accessed 4 times a year – S3 infrequently access
- Data preserved- S3 Glacier Deep Archive

upvoted 9 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #374

Topic 1

A company is running several business applications in three separate VPCs within the us-east-1 Region. The applications must be able to communicate between VPCs. The applications also must be able to consistently send hundreds of gigabytes of data each day to a latency-sensitive application that runs in a single on-premises data center.

A solutions architect needs to design a network connectivity solution that maximizes cost-effectiveness.

Which solution meets these requirements?

- A. Configure three AWS Site-to-Site VPN connections from the data center to AWS. Establish connectivity by configuring one VPN connection for each VPC.
- B. Launch a third-party virtual network appliance in each VPC. Establish an IPsec VPN tunnel between the data center and each virtual appliance.
- C. Set up three AWS Direct Connect connections from the data center to a Direct Connect gateway in us-east-1. Establish connectivity by configuring each VPC to use one of the Direct Connect connections.
- D. Set up one AWS Direct Connect connection from the data center to AWS. Create a transit gateway, and attach each VPC to the transit gateway. Establish connectivity between the Direct Connect connection and the transit gateway. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

TariqKipkemei **Highly Voted** 1 year, 1 month ago

Selected Answer: D

AWS Transit Gateway connects your Amazon Virtual Private Clouds (VPCs) and on-premises networks through a central hub. This connection simplifies your network and puts an end to complex peering relationships. Transit Gateway acts as a highly scalable cloud router—each new connection is made only once.

<https://aws.amazon.com/transit-gateway/#:~:text=AWS-,Transit%20Gateway,-connects%20your%20Amazon>

upvoted 6 times

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: D

Selected Answer: D

This option leverages a single Direct Connect for consistent, private connectivity between the data center and AWS. The transit gateway allows each VPC to share the Direct Connect while keeping the VPCs isolated. This provides a cost-effective architecture to meet the requirements.

upvoted 5 times

upliftinghut Most Recent 10 months, 2 weeks ago

Selected Answer: D

AWS Direct connect is costly but the saving comes from less data transfer cost with Direct Connect and Transit gateway

upvoted 4 times

alexandercamachop 1 year, 6 months ago

Selected Answer: D

Transit GW, is a hub for connecting all VPCs.

Direct Connect is expensive, therefor only 1 of them connected to the Transit GW (Hub for all our VPCs that we connect to it)

upvoted 4 times

KMohsoe 1 year, 6 months ago

Selected Answer: D

Option D

upvoted 4 times

Sivasaa 1 year, 7 months ago

Can someone tell why option C will not work here

upvoted 3 times

pentium75 11 months, 2 weeks ago

And besides the cost, C does not allow the applications "to communicate between VPCs".

upvoted 2 times

jdamaian 1 year, 7 months ago

cost-effectiveness, 3 DC are more than 1 (more expensive). There is no need to connect more than 1 DC.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Using multiple Site-to-Site VPNs (A) or Direct Connects (C) incurs higher costs without providing significant benefits.

upvoted 2 times

SkyZeroZx 1 year, 7 months ago

Selected Answer: D

cost-effectiveness

D

upvoted 2 times

WhericanIstart 1 year, 8 months ago

Selected Answer: D

Transit Gateway will achieve this result..

upvoted 4 times

Karlos99 1 year, 8 months ago

Selected Answer: D

maximizes cost-effectiveness

upvoted 3 times

[Removed] 1 year, 9 months ago

Selected Answer: D

ddddddddd
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #375

Topic 1

An ecommerce company is building a distributed application that involves several serverless functions and AWS services to complete order-processing tasks. These tasks require manual approvals as part of the workflow. A solutions architect needs to design an architecture for the order-processing application. The solution must be able to combine multiple AWS Lambda functions into responsive serverless applications. The solution also must orchestrate data and services that run on Amazon EC2 instances, containers, or on-premises servers.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use AWS Step Functions to build the application. **Most Voted**
- B. Integrate all the application components in an AWS Glue job.
- C. Use Amazon Simple Queue Service (Amazon SQS) to build the application.
- D. Use AWS Lambda functions and Amazon EventBridge events to build the application.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**kinglong12** **Highly Voted** 1 year, 2 months ago**Selected Answer: A**

AWS Step Functions is a fully managed service that makes it easy to build applications by coordinating the components of distributed applications and microservices using visual workflows. With Step Functions, you can combine multiple AWS Lambda functions into responsive serverless applications and orchestrate data and services that run on Amazon EC2 instances, containers, or on-premises servers. Step Functions also allows for manual approvals as part of the workflow. This solution meets all the requirements with the least operational overhead.

upvoted 15 times

Guru4Cloud **Highly Voted** 9 months, 1 week ago**Selected Answer: A**

AWS Step Functions allow you to easily coordinate multiple Lambda functions and services into serverless workflows with visual workflows. Step Functions are designed for building distributed applications that combine services and require human approval steps.

Using Step Functions provides a fully managed orchestration service with minimal operational overhead.
upvoted 7 times

TariqKipkemei Most Recent 7 months, 3 weeks ago

Selected Answer: A

involves several serverless functions and AWS services, require manual approvals as part of the workflow, combine the Lambda functions into responsive serverless applications, orchestrate data and services = AWS Step Functions
upvoted 4 times

capino 9 months, 3 weeks ago

Selected Answer: A

Serverless workflow service that need human approval:::step functions
upvoted 3 times

BeeKayEnn 1 year, 2 months ago

Key: Distributed Application Processing, Microservices orchestration (Orchestrate Data and Services)

A would be the best fit.

AWS Step Functions is a visual workflow service that helps developers use AWS services to build distributed applications, automate processes, orchestrate microservices, and create data and machine learning (ML) pipelines.

Reference: [https://aws.amazon.com/step-functions/#:~:text=AWS%20Step%20Functions%20is%20a,machine%20learning%20\(ML\)%20pipelines](https://aws.amazon.com/step-functions/#:~:text=AWS%20Step%20Functions%20is%20a,machine%20learning%20(ML)%20pipelines).

upvoted 4 times

COTIT 1 year, 2 months ago

Selected Answer: A

Approval is explicit for the solution. -> "A common use case for AWS Step Functions is a task that requires human intervention (for example, an approval process). Step Functions makes it easy to coordinate the components of distributed applications as a series of steps in a visual workflow called a state machine. You can quickly build and run state machines to execute the steps of your application in a reliable and scalable fashion. (<https://aws.amazon.com/pt/blogs/compute/implementing-serverless-manual-approval-steps-in-aws-step-functions-and-amazon-api-gateway/>)"

upvoted 6 times

ktulu2602 1 year, 2 months ago

Selected Answer: A

Option A: Use AWS Step Functions to build the application.

AWS Step Functions is a serverless workflow service that makes it easy to coordinate distributed applications and microservices using visual workflows. It is an ideal solution for designing architectures for distributed applications that involve multiple AWS services and serverless functions, as it allows us to orchestrate the flow of our application components using visual workflows. AWS Step Functions also integrates with other AWS services like AWS Lambda, Amazon EC2, and Amazon ECS, and it has built-in error handling and retry mechanisms. This option provides a serverless solution with the least operational overhead for building the application.

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #376

Topic 1

A company has launched an Amazon RDS for MySQL DB instance. Most of the connections to the database come from serverless applications. Application traffic to the database changes significantly at random intervals. At times of high demand, users report that their applications experience database connection rejection errors.

Which solution will resolve this issue with the LEAST operational overhead?

- A. Create a proxy in RDS Proxy. Configure the users' applications to use the DB instance through RDS Proxy. **Most Voted**
- B. Deploy Amazon ElastiCache for Memcached between the users' applications and the DB instance.
- C. Migrate the DB instance to a different instance class that has higher I/O capacity. Configure the users' applications to use the new DB instance.
- D. Configure Multi-AZ for the DB instance. Configure the users' applications to switch between the DB instances.

Correct Answer: A

Community vote distribution

A (100%)

Comments

TariqKipkemei **Highly Voted** 7 months, 3 weeks ago

Selected Answer: A

database connection rejection errors = RDS Proxy
upvoted 7 times

ktulu2602 **Highly Voted** 1 year, 2 months ago

Selected Answer: A

The correct solution for this scenario would be to create a proxy in RDS Proxy. RDS Proxy allows for managing thousands of concurrent database connections, which can help reduce connection errors. RDS Proxy also provides features such as connection pooling, read/write splitting, and retries. This solution requires the least operational overhead as it does not involve migrating to a different instance class or setting up a new cache layer. Therefore, option A is the correct answer.
upvoted 5 times

Guru4Cloud **Most Recent** 9 months, 1 week ago

Selected Answer: A

RDS Proxy provides a proxy layer that pools and shares database connections to improve scalability. This allows the proxy to handle connection spikes to the database gracefully.

Using RDS Proxy requires minimal operational overhead - just create the proxy and reconfigure applications to use it. No code changes needed.

upvoted 4 times

antropaws 1 year ago

Wait, why not B?????

upvoted 2 times

live_reply_developers 11 months, 1 week ago

Amazon ElastiCache tends to have a lower operational overhead compared to Amazon RDS Proxy. BUT we already have "Amazon RDS for MySQL DB instance"

upvoted 2 times

Guru4Cloud 9 months, 1 week ago

ElastiCache (B) and larger instance type (C) help performance but don't resolve connection issues.

upvoted 2 times

Guru4Cloud 9 months, 1 week ago

ElastiCache (B) and larger instance type (C) help performance but don't resolve connection issues.

upvoted 4 times

roxx529 1 year ago

To reduce application failures resulting from database connection timeouts, the best solution is to enable RDS Proxy on the RDS DB instances

upvoted 2 times

COTIT 1 year, 2 months ago

Selected Answer: A

Many applications, including those built on modern serverless architectures, can have a large number of open connections to the database server and may open and close database connections at a high rate, exhausting database memory and compute resources. Amazon RDS Proxy allows applications to pool and share connections established with the database, improving database efficiency and application scalability. (<https://aws.amazon.com/pt/rds/proxy/>)

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #377

Topic 1

A company recently deployed a new auditing system to centralize information about operating system versions, patching, and installed software for Amazon EC2 instances. A solutions architect must ensure all instances provisioned through EC2 Auto Scaling groups successfully send reports to the auditing system as soon as they are launched and terminated.

Which solution achieves these goals MOST efficiently?

- A. Use a scheduled AWS Lambda function and run a script remotely on all EC2 instances to send data to the audit system.
- B. Use EC2 Auto Scaling lifecycle hooks to run a custom script to send data to the audit system when instances are launched and terminated. **Most Voted**
- C. Use an EC2 Auto Scaling launch configuration to run a custom script through user data to send data to the audit system when instances are launched and terminated.
- D. Run a custom script on the instance operating system to send data to the audit system. Configure the script to be invoked by the EC2 Auto Scaling group when the instance starts and is terminated.

Correct Answer: B*Community vote distribution*B (100%)**Comments****ktulu2602** **Highly Voted** 2 years, 2 months ago**Selected Answer: B**

The most efficient solution for this scenario is to use EC2 Auto Scaling lifecycle hooks to run a custom script to send data to the audit system when instances are launched and terminated. The lifecycle hook can be used to delay instance termination until the script has completed, ensuring that all data is sent to the audit system before the instance is terminated. This solution is more efficient than using a scheduled AWS Lambda function, which would require running the function periodically and may not capture all instances launched and terminated within the interval. Running a custom script through user data is also not an optimal solution, as it may not guarantee that all instances send data to the audit system. Therefore, option B is the correct answer.

upvoted 12 times

COTIT **Highly Voted** 2 years, 2 months ago**Selected Answer: B**

Amazon EC2 Auto Scaling offers the ability to add lifecycle hooks to your Auto Scaling groups. These hooks let you create

Amazon EC2 Auto Scaling offers the ability to add lifecycle hooks to your Auto Scaling groups. These hooks let you create solutions that are aware of events in the Auto Scaling instance lifecycle, and then perform a custom action on instances when the corresponding lifecycle event occurs. (<https://docs.aws.amazon.com/autoscaling/ec2/userguide/lifecycle-hooks.html>)

upvoted 6 times

LeonSauveterre Most Recent 6 months, 1 week ago

Selected Answer: B

C: a ****launch**** configuration CANNOT handle instance termination automatically.

upvoted 1 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: B

Use EC2 Auto Scaling lifecycle hooks to run a custom script to send data to the audit system when instances are launched and terminated

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

EC2 Auto Scaling lifecycle hooks allow you to perform custom actions as instances launch and terminate. This is the most efficient way to trigger the auditing script execution at instance launch and termination.

upvoted 5 times

WheracanIstart 2 years, 2 months ago

Selected Answer: B

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/lifecycle-hooks.html>

upvoted 4 times

fkie4 2 years, 2 months ago

it is B. read this:

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/lifecycle-hooks.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #378

Topic 1

A company is developing a real-time multiplayer game that uses UDP for communications between the client and servers in an Auto Scaling group. Spikes in demand are anticipated during the day, so the game server platform must adapt accordingly. Developers want to store gamer scores and other non-relational data in a database solution that will scale without intervention.

Which solution should a solutions architect recommend?

- A. Use Amazon Route 53 for traffic distribution and Amazon Aurora Serverless for data storage.
- B. Use a Network Load Balancer for traffic distribution and Amazon DynamoDB on-demand for data storage. **Most Voted**
- C. Use a Network Load Balancer for traffic distribution and Amazon Aurora Global Database for data storage.
- D. Use an Application Load Balancer for traffic distribution and Amazon DynamoDB global tables for data storage.

Correct Answer: B

Community vote distribution

B (100%)

Comments

TariqKipkemei **Highly Voted** 1 year ago

Selected Answer: B

UDP = NLB
Non-relational data = Dynamo DB
upvoted 17 times

fruto123 **Highly Voted** 1 year, 2 months ago

Selected Answer: B

key words - UDP, non-relational data
answers - NLB for UDP application, DynamoDB for non-relational data
upvoted 6 times

Manimgh **Most Recent** 1 month, 3 weeks ago

Selected Answer: B

non-relational data = DynamoDB

upvoted 1 times

Guru4Cloud 9 months, 1 week ago

Selected Answer: B

This option provides the most scalable and optimized architecture for the real-time multiplayer game:

Network Load Balancer efficiently distributes UDP gaming traffic to the Auto Scaling group of game servers.

DynamoDB On-Demand mode provides auto-scaling non-relational data storage for gamer scores and other game data.

DynamoDB is optimized for fast, high-scale access patterns seen in gaming.

Together, the Network Load Balancer and DynamoDB On-Demand provide an architecture that can smoothly scale up and down to match spikes in gaming demand.

upvoted 5 times

elarningtakai 1 year, 2 months ago

Selected Answer: B

Option B is a good fit because a Network Load Balancer can handle UDP traffic, and Amazon DynamoDB on-demand can provide automatic scaling without intervention

upvoted 3 times

KAUS2 1 year, 2 months ago

Selected Answer: B

Correct option is "B"

upvoted 2 times

aragon_saa 1 year, 2 months ago

B

<https://www.examttopics.com/discussions/amazon/view/29756-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Kenp1192 1 year, 2 months ago

B

Because NLB can handle UDP and DynamoDB is Non-Relational

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #379

Topic 1

A company hosts a frontend application that uses an Amazon API Gateway API backend that is integrated with AWS Lambda. When the API receives requests, the Lambda function loads many libraries. Then the Lambda function connects to an Amazon RDS database, processes the data, and returns the data to the frontend application. The company wants to ensure that response latency is as low as possible for all its users with the fewest number of changes to the company's operations.

Which solution will meet these requirements?

- A. Establish a connection between the frontend application and the database to make queries faster by bypassing the API.
- B. Configure provisioned concurrency for the Lambda function that handles the requests. **Most Voted**
- C. Cache the results of the queries in Amazon S3 for faster retrieval of similar datasets.
- D. Increase the size of the database to increase the number of connections Lambda can establish at one time.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**UnluckyDucky** **Highly Voted** 1 year, 2 months ago**Selected Answer: B**

Key: the Lambda function loads many libraries

Configuring provisioned concurrency would get rid of the "cold start" of the function therefore speeding up the process.
upvoted 17 times

kampatra **Highly Voted** 1 year, 2 months ago**Selected Answer: B**

Provisioned concurrency – Provisioned concurrency initializes a requested number of execution environments so that they are prepared to respond immediately to your function's invocations. Note that configuring provisioned concurrency incurs charges to your AWS account.

upvoted 11 times

TariqKipkemei **Most Recent** 7 months, 3 weeks ago**Selected Answer: B**

Provisioned concurrency pre-initializes execution environments which are prepared to respond immediately to incoming function requests.

upvoted 7 times

Guru4Cloud 9 months, 1 week ago

Selected Answer: B

Provisioned concurrency ensures a configured number of execution environments are ready to serve requests to the Lambda function. This avoids cold starts where the function would otherwise need to load all the libraries on each invocation.

upvoted 4 times

Guru4Cloud 9 months, 1 week ago

Selected Answer: B

Provisioned concurrency ensures a configured number of execution environments are ready to serve requests to the Lambda function. This avoids cold starts where the function would otherwise need to load all the libraries on each invocation.

upvoted 2 times

elearningtakai 1 year, 2 months ago

Selected Answer: B

Answer B is correct

<https://docs.aws.amazon.com/lambda/latest/dg/provisioned-concurrency.html>

Answer C: need to modify the application

upvoted 5 times

elearningtakai 1 year, 2 months ago

This is relevant to "cold start" with keywords: "Lambda function loads many libraries"

upvoted 2 times

Karlos99 1 year, 2 months ago

Selected Answer: B

<https://docs.aws.amazon.com/lambda/latest/dg/provisioned-concurrency.html>

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #380

Topic 1

A company is migrating its on-premises workload to the AWS Cloud. The company already uses several Amazon EC2 instances and Amazon RDS DB instances. The company wants a solution that automatically starts and stops the EC2 instances and DB instances outside of business hours. The solution must minimize cost and infrastructure maintenance.

Which solution will meet these requirements?

- A. Scale the EC2 instances by using elastic resize. Scale the DB instances to zero outside of business hours.
- B. Explore AWS Marketplace for partner solutions that will automatically start and stop the EC2 instances and DB instances on a schedule.
- C. Launch another EC2 instance. Configure a crontab schedule to run shell scripts that will start and stop the existing EC2 instances and DB instances on a schedule.
- D. Create an AWS Lambda function that will start and stop the EC2 instances and DB instances. Configure Amazon EventBridge to invoke the Lambda function on a schedule. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

ktulu2602 **Highly Voted** 2 years, 2 months ago

Selected Answer: D

The most efficient solution for automatically starting and stopping EC2 instances and DB instances on a schedule while minimizing cost and infrastructure maintenance is to create an AWS Lambda function and configure Amazon EventBridge to invoke the function on a schedule.

Option A, scaling EC2 instances by using elastic resize and scaling DB instances to zero outside of business hours, is not feasible as DB instances cannot be scaled to zero.

Option B, exploring AWS Marketplace for partner solutions, may be an option, but it may not be the most efficient solution and could potentially add additional costs.

Option C, launching another EC2 instance and configuring a crontab schedule to run shell scripts that will start and stop the existing EC2 instances and DB instances on a schedule, adds unnecessary infrastructure and maintenance.

upvoted 17 times

Guru4Cloud Highly Voted 1 year, 9 months ago

Selected Answer: D

This option leverages AWS Lambda and EventBridge to automatically schedule the starting and stopping of resources.

Lambda provides the script/code to stop/start instances without managing servers.

EventBridge triggers the Lambda on a schedule without cronjobs.

No additional code or third party tools needed.

Serverless, maintenance-free solution

upvoted 6 times

1e22522 Most Recent 10 months, 1 week ago

Selected Answer: D

its d but nowadays u use system manager me thinks

upvoted 3 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: D

Create an AWS Lambda function that will start and stop the EC2 instances and DB instances. Configure Amazon EventBridge to invoke the Lambda function on a schedule.

upvoted 4 times

WheracanIstart 2 years, 2 months ago

Selected Answer: D

Minimize cost and maintenance...

upvoted 3 times

[Removed] 2 years, 2 months ago

Selected Answer: D

DDDDDDDDDDDD

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #381


Topic 1

A company hosts a three-tier web application that includes a PostgreSQL database. The database stores the metadata from documents. The company searches the metadata for key terms to retrieve documents that the company reviews in a report each month. The documents are stored in Amazon S3. The documents are usually written only once, but they are updated frequently.

The reporting process takes a few hours with the use of relational queries. The reporting process must not prevent any document modifications or the addition of new documents. A solutions architect needs to implement a solution to speed up the reporting process.

Which solution will meet these requirements with the LEAST amount of change to the application code?

- A. Set up a new Amazon DocumentDB (with MongoDB compatibility) cluster that includes a read replica. Scale the read replica to generate the reports.
- B. Set up a new Amazon Aurora PostgreSQL DB cluster that includes an Aurora Replica. Issue queries to the Aurora Replica to generate the reports. **Most Voted**
- C. Set up a new Amazon RDS for PostgreSQL Multi-AZ DB instance. Configure the reporting module to query the secondary RDS node so that the reporting module does not affect the primary node.
- D. Set up a new Amazon DynamoDB table to store the documents. Use a fixed write capacity to support new document entries. Automatically scale the read capacity to support the reports.

Correct Answer: B*Community vote distribution***Comments****Guru4Cloud** **Highly Voted** 1 year, 8 months ago**Selected Answer: B**

The key reasons are:

Aurora PostgreSQL provides native PostgreSQL compatibility, so minimal code changes would be required. Using an Aurora Replica separates the reporting workload from the main workload, preventing any slowdown of document updates/inserts.

updates/inserts.

Aurora can auto-scale read replicas to handle the reporting load.

This allows leveraging the existing PostgreSQL database without major changes. DynamoDB would require more significant rewrite of data access code.

RDS Multi-AZ alone would not fully separate the workloads, as the secondary is for HA/failover more than scaling read workloads.

upvoted 12 times

TariqKipkemei Highly Voted 2 years ago

Selected Answer: B

Load balancing = Read replica

High availability = Multi AZ

upvoted 8 times

BillaRanga 1 year, 4 months ago

No Modifications allowed = Read Replica

upvoted 4 times

terminator69 Most Recent 9 months, 3 weeks ago

How in the bloody hell it's D?????

upvoted 1 times

TruthWS 1 year, 2 months ago

B is correct

upvoted 2 times

ExamGuru727 1 year, 2 months ago

Selected Answer: B

We also have a requirement for the Least amount of change to the code.

Since our DB is PostgreSQL, A & D are immediately out.

Multi-AZ won't help with offloading read requests, hence the answer is B ;)

upvoted 5 times

Buck12345 1 year, 3 months ago

It is B

upvoted 2 times

Cyberkayu 1 year, 5 months ago

Selected Answer: C

D. Reporting process Must not prevent = allow modification and addition of new document.

all read replica were wrong.

upvoted 1 times

pentium75 1 year, 5 months ago

How would 'issuing queries to the read replica' prevent modifications or updates?

upvoted 2 times

KMohsoe 2 years ago

Selected Answer: A

Why not A? :(

upvoted 1 times

wRhIH 1 year, 11 months ago

"The reporting process takes a few hours with the use of RELATIONAL queries."

upvoted 4 times

Murtadhaceit 1 year, 6 months ago

DocumentDB (For MongoDB) is no SQL. DynamoDB is also No SQL. Therefore, options A and D are out.
upvoted 5 times

lexotan 2 years, 1 month ago

Selected Answer: B

B is the right one. why admin does not correct these wrong answers?
upvoted 4 times

imvb88 2 years, 1 month ago

Selected Answer: B

The reporting process queries the metadata (not the documents) and use relational queries-> A, D out
C: wrong since secondary RDS node in MultiAZ setup is in standby mode, not available for querying
B: reporting using a Replica is a design pattern. Using Aurora is an exam pattern.
upvoted 5 times

WheracanIstart 2 years, 2 months ago

Selected Answer: B

B is right..
upvoted 2 times

Maximus007 2 years, 2 months ago

Selected Answer: B

While both B&D seems to be a relevant, ChatGPT suggest B as a correct one
upvoted 2 times

cegama543 2 years, 2 months ago

Selected Answer: B

Option B (Set up a new Amazon Aurora PostgreSQL DB cluster that includes an Aurora Replica. Issue queries to the Aurora Replica to generate the reports) is the best option for speeding up the reporting process for a three-tier web application that includes a PostgreSQL database storing metadata from documents, while not impacting document modifications or additions, with the least amount of change to the application code.
upvoted 3 times

UnluckyDucky 2 years, 2 months ago

Selected Answer: B

"LEAST amount of change to the application code"

Aurora is a relational database, it supports PostgreSQL and with the help of read replicas we can issue the reporting process that take several hours to the replica, therefore not affecting the primary node which can handle new writes or document modifications.
upvoted 2 times

Ashukaushal619 2 years, 2 months ago

its D only ,recorrected
upvoted 2 times

Murtadhaceit 1 year, 6 months ago

DynamoDB is no SQL. A and D are out!
upvoted 2 times

Ashukaushal619 2 years, 2 months ago

Selected Answer: B

bbbbbbbbb
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #382

Topic 1

A company has a three-tier application on AWS that ingests sensor data from its users' devices. The traffic flows through a Network Load Balancer (NLB), then to Amazon EC2 instances for the web tier, and finally to EC2 instances for the application tier. The application tier makes calls to a database.

What should a solutions architect do to improve the security of the data in transit?

- A. Configure a TLS listener. Deploy the server certificate on the NLB. **Most Voted**
- B. Configure AWS Shield Advanced. Enable AWS WAF on the NLB.
- C. Change the load balancer to an Application Load Balancer (ALB). Enable AWS WAF on the ALB.
- D. Encrypt the Amazon Elastic Block Store (Amazon EBS) volume on the EC2 instances by using AWS Key Management Service (AWS KMS).

Correct Answer: A

Community vote distribution

A (100%)

Comments

fruto123 **Highly Voted** 1 year, 2 months ago

Selected Answer: A

Network Load Balancers now support TLS protocol. With this launch, you can now offload resource intensive decryption/encryption from your application servers to a high throughput, and low latency Network Load Balancer. Network Load Balancer is now able to terminate TLS traffic and set up connections with your targets either over TCP or TLS protocol.

<https://docs.aws.amazon.com/elasticloadbalancing/latest/network/create-tls-listener.html>

https://exampleloadbalancer.com/nlbtl_demo.html

upvoted 23 times

imvb88 **Highly Voted** 1 year, 1 month ago

Selected Answer: A

security of data in transit -> think of SSL/TLS. Check: NLB supports TLS

<https://docs.aws.amazon.com/elasticloadbalancing/latest/network/create-tls-listener.html>

B (DDoS) C (SQL Injection) D (EBS) is for data at rest

B (DDOS), C (SQL Injection), D (EBS) is for data at rest.

upvoted 18 times

TariqKipkemei Most Recent 7 months, 3 weeks ago

Selected Answer: A

secure data in transit = TLS

upvoted 2 times

Guru4Cloud 9 months, 1 week ago

Selected Answer: A

TLS provides encryption for data in motion over the network, protecting against eavesdropping and tampering. A valid server certificate signed by a trusted CA will provide further security.

upvoted 6 times

klayytech 1 year, 2 months ago

Selected Answer: A

To improve the security of data in transit, you can configure a TLS listener on the Network Load Balancer (NLB) and deploy the server certificate on it. This will encrypt traffic between clients and the NLB. You can also use AWS Certificate Manager (ACM) to provision, manage, and deploy SSL/TLS certificates for use with AWS services and your internal connected resources¹.

You can also change the load balancer to an Application Load Balancer (ALB) and enable AWS WAF on it. AWS WAF is a web application firewall that helps protect your web applications from common web exploits that could affect application availability, compromise security, or consume excessive resources³.

the A and C correct without transit but the need to improve the security of the data in transit? so he need SSL/TLS certificates

upvoted 3 times

Maximus007 1 year, 2 months ago

Selected Answer: A

agree with fruto123

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #383

Topic 1

A company is planning to migrate a commercial off-the-shelf application from its on-premises data center to AWS. The software has a software licensing model using sockets and cores with predictable capacity and uptime requirements. The company wants to use its existing licenses, which were purchased earlier this year.

Which Amazon EC2 pricing option is the MOST cost-effective?

A. Dedicated Reserved Hosts **Most Voted**


B. Dedicated On-Demand Hosts

C. Dedicated Reserved Instances

D. Dedicated On-Demand Instances

Correct Answer: A

Community vote distribution

**Comments**

fkie4 **Highly Voted** 2 years, 2 months ago

Selected Answer: A

"predictable capacity and uptime requirements" means "Reserved"

"sockets and cores" means "dedicated host"

upvoted 21 times

imvb88 **Highly Voted** 2 years, 1 month ago

Selected Answer: A

Bring custom purchased licenses to AWS -> Dedicated Host -> C,D out

Need cost effective solution -> "reserved" -> A

upvoted 7 times

imvb88 2 years, 1 month ago

<https://aws.amazon.com/ec2/dedicated-hosts/>

Amazon EC2 Dedicated Hosts allow you to use your eligible software licenses from vendors such as Microsoft and Oracle on Amazon EC2, so that you get the flexibility and cost effectiveness of using your own licenses, but with the simplicity

Amazon EC2, so that you get the flexibility and cost effectiveness of using your own licenses, but with the resiliency, simplicity and elasticity of AWS.

upvoted 3 times

Hrishi_707 Most Recent 1 year, 3 months ago

BYOL >>> Dedicated Hosts

upvoted 4 times

Uzbekistan 1 year, 3 months ago

Selected Answer: A

A. Dedicated Reserved Hosts

Here's why:

License Flexibility: Dedicated Reserved Hosts allow the company to bring their existing licenses to AWS. This option enables them to continue using their purchased licenses without any additional cost or licensing changes.

Cost Optimization: Reserved Hosts offer significant cost savings compared to On-Demand pricing. By purchasing Reserved Hosts, the company can benefit from discounted hourly rates for the entire term of the reservation, which typically spans one or three years.

upvoted 3 times

jjcode 1 year, 3 months ago

I work with COTS applications they require a three tier architecture, its completely irrelevant and confusing to add that to the question, the key word here is licenses, since AWS wants your to use their solutions the answer to this is which of one the options solves this particular problem, in this case its dedicated hosts.

upvoted 1 times

BillaRanga 1 year, 4 months ago

Selected Answer: A

What is difference between dedicated host and reserved instance?

Dedicated Instance: The physical machine or underlying hardware is reserved for use for the whole account. You can have instances for different purposes on this hardware. Dedicated Host: The physical machine or the underlying hardware is reserved for "Single Use" only, eg. a certain application.

upvoted 2 times

BillaRanga 1 year, 4 months ago

What is the difference between a dedicated instance and a dedicated host tenancy?

Dedicated Instance (dedicated) — Your instance runs on single-tenant hardware. Dedicated Host (host) — Your instance runs on a physical server with EC2 instance capacity fully dedicated to your use, an isolated server with configurations that you can control.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: A

Actually the question is a bit ambiguous because there ARE "software licensing model using sockets and cores" that accept virtual sockets are cores as the base, for which C would work. But most of these license models are based on PHYSICAL sockets, thus A.

upvoted 4 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: A

Dedicated Hosts give you visibility and control over how instances are placed on a physical server and also enable you to use your existing server-bound software licenses like Windows Server

upvoted 3 times

wsdasdasdqwda 1 year, 7 months ago

Easy with one, but only 79% up to now answered correctly. It is A. Reserved because of the predictable and sockets and cores means dedicated host.

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

The correct answer is C. Dedicated Reserved Instances.

Dedicated Reserved Instances (DRIs) are the most cost-effective option for workloads that have predictable capacity and uptime requirements. DRIs offer a significant discount over On-Demand Instances, and they can be used to lock in a price for a period of time.

In this case, the company has predictable capacity and uptime requirements because the software has a software licensing model using sockets and cores. The company also wants to use its existing licenses, which were purchased earlier this year. Therefore, DRIs are the most cost-effective option.

upvoted 3 times

riccardoto 1 year, 10 months ago

Selected Answer: C

I don't agree with people voting "A". The question reference that the COTS Application has a licensing model based on "sockets and cores". The question does not specify if it means TCP sockets (= open connections) or hardware sockets, so I assume that "TCP sockets are intended". If this is the case, sockets and cores can also remain stable with reserved instances - which are cheaper than reserved hosts.

I would go with "A" only if the question would clearly state that the COTS application has some strong dependency on physical hardware.

upvoted 1 times

JA2018 6 months, 2 weeks ago

Hi Riccardoto, perhaps you might to enlighten us which software licensing model will bill customer on a "per-TCP socket" basis? :-)

Cheers.

upvoted 1 times

riccardoto 1 year, 10 months ago

note: instead, if by socket we mean "CPU sockets", then A would be the right one.

upvoted 2 times

pentium75 1 year, 5 months ago

Even if "sockets" mean TCP sockets there are still the cores, thus A

upvoted 3 times

aragon_saa 2 years, 2 months ago

A

<https://www.examttopics.com/discussions/amazon/view/35818-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

fruto123 2 years, 2 months ago

Selected Answer: A

Dedicated Host Reservations provide a billing discount compared to running On-Demand Dedicated Hosts. Reservations are available in three payment options.

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/dedicated-hosts-overview.html>

upvoted 4 times

Kenp1192 2 years, 2 months ago

A

is the most cost effective

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #384

Topic 1

A company runs an application on Amazon EC2 Linux instances across multiple Availability Zones. The application needs a storage layer that is highly available and Portable Operating System Interface (POSIX)-compliant. The storage layer must provide maximum data durability and must be shareable across the EC2 instances. The data in the storage layer will be accessed frequently for the first 30 days and will be accessed infrequently after that time.

Which solution will meet these requirements MOST cost-effectively?

- A. Use the Amazon S3 Standard storage class. Create an S3 Lifecycle policy to move infrequently accessed data to S3 Glacier.
- B. Use the Amazon S3 Standard storage class. Create an S3 Lifecycle policy to move infrequently accessed data to S3 Standard-Infrequent Access (S3 Standard-IA).
- C. Use the Amazon Elastic File System (Amazon EFS) Standard storage class. Create a lifecycle management policy to move infrequently accessed data to EFS Standard-Infrequent Access (EFS Standard-IA). **Most Voted**
- D. Use the Amazon Elastic File System (Amazon EFS) One Zone storage class. Create a lifecycle management policy to move infrequently accessed data to EFS One Zone-Infrequent Access (EFS One Zone-IA).

Correct Answer: C*Community vote distribution*

C (93%)

D (7%)

Comments**TariqKipkemei** **Highly Voted** 2 years ago**Selected Answer: C**

Multi AZ = both EFS and S3 support
Storage classes = both EFS and S3 support
POSIX file system access = only Amazon EFS supports
upvoted 17 times

LazyTs **Highly Voted** 1 year, 9 months ago**Selected Answer: C**

POSIX => EFS
<https://docs.aws.amazon.com/efs/latest/ua/whatisefs.html>

upvoted 8 times

Danilus Most Recent 6 months, 4 weeks ago

Selected Answer: C

key-Portable Operating System Interface (POSIX)-compliant.
key -shareable across the ec2 instances
key-accessed frequently for the first 30 days and will be accessed infrequently after
key-storage layer that is highly available

so the only service that work with POSIX is EFS in this case
s3 can not use posix so A and B are wrong answers
is not D because is not highly available because is One Zone so the answer is c
upvoted 2 times

zinabu 1 year, 2 months ago

Answer:B
cause there is no life cycle policy for EFS that will work in S3 only.
upvoted 2 times

jjcode 1 year, 3 months ago

"storage layer will be accessed frequently for the first 30 days and will be accessed infrequently after that time" Was the only reason they added this to trick you?
upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: C

POSIX -> EFS, "maximum data durability" rules out One Zone
upvoted 4 times

maudsha 1 year, 7 months ago

Selected Answer: C

Both standard and one zone have same durability.
<https://docs.aws.amazon.com/efs/latest/ug/storage-classes.html>

Also EFS one zone can work with multiple EC2s in different AZs. But there will be a cost involved when you are accessing the EFS from a different AZ EC2. (EC2 data access charges)
<https://docs.aws.amazon.com/efs/latest/ug/how-it-works.html>
So if "all" EC2 instances accessing the files frequently there will be a storage cost + EC2 data access charges if you choose one zone.

So i would choose C.
upvoted 2 times

beast2091 1 year, 7 months ago

Ans: C
upvoted 2 times

baba365 1 year, 8 months ago

Ans: D, one-zone IA for 'most cost effective' .

<https://aws.amazon.com/efs/features/infrequent-access/>
upvoted 1 times

AAAWrekng 1 year, 7 months ago

How does D fulfill the data durability requirement? Requirements must be met first, then consider 'most cost effective' - if you go to a tire shop, and say you want 4 new tires as cheap as possible. And they take off 4 tires and put on 2... Then they say you wanted it as cheap as possible...
upvoted 4 times

Gajendr 1 year, 5 months ago

What about " The application needs a storage layer that is highly available" and "application on Amazon EC2 Linux instances

What about "The application needs a storage layer that is highly available" and "application on Amazon EC2 Linux instances across multiple Availability Zones" ?

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

Use the Amazon Elastic File System (Amazon EFS) Standard storage class. Create a lifecycle management policy to move infrequently accessed data to EFS Standard-Infrequent Access (EFS Standard-IA).

upvoted 3 times

[Removed] 1 year, 11 months ago

Selected Answer: D

Amazon Elastic File System (Amazon EFS) Standard storage class = "maximum data durability"

upvoted 2 times

pentium75 1 year, 5 months ago

"ONE ZONE-IA" does not meet the "maximum data durability" requirement

upvoted 2 times

Yadav_Sanjay 1 year, 11 months ago

Selected Answer: D

D - It should be cost-effective

upvoted 2 times

pentium75 1 year, 5 months ago

But D does meet the durability requirement.

upvoted 1 times

Abrar2022 1 year, 12 months ago

Selected Answer: C

POSIX file system access = only Amazon EFS supports

upvoted 4 times

imvb88 2 years, 1 month ago

Selected Answer: C

POSIX + sharable across EC2 instances --> EFS --> A, B out

Instances run across multiple AZ -> C is needed.

upvoted 2 times

WherecanIstart 2 years, 2 months ago

Selected Answer: C

Linux based system points to EFS plus POSIX-compliant is also EFS related.

upvoted 3 times

fk1e4 2 years, 2 months ago

Selected Answer: C

"POSIX-compliant" means EFS.

also, file system can be shared with multiple EC2 instances means "EFS"

upvoted 5 times

KAUS2 2 years, 2 months ago

Selected Answer: C

Option C is the correct answer .

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #385

Topic 1

A solutions architect is creating a new VPC design. There are two public subnets for the load balancer, two private subnets for web servers, and two private subnets for MySQL. The web servers use only HTTPS. The solutions architect has already created a security group for the load balancer allowing port 443 from 0.0.0.0/0. Company policy requires that each resource has the least access required to still be able to perform its tasks.

Which additional configuration strategy should the solutions architect use to meet these requirements?

- A. Create a security group for the web servers and allow port 443 from 0.0.0.0/0. Create a security group for the MySQL servers and allow port 3306 from the web servers security group.
- B. Create a network ACL for the web servers and allow port 443 from 0.0.0.0/0. Create a network ACL for the MySQL servers and allow port 3306 from the web servers security group.
- C. Create a security group for the web servers and allow port 443 from the load balancer. Create a security group for the MySQL servers and allow port 3306 from the web servers security group. **Most Voted**
- D. Create a network ACL for the web servers and allow port 443 from the load balancer. Create a network ACL for the MySQL servers and allow port 3306 from the web servers security group.

Correct Answer: C

Community vote distribution

C (100%)

Comments

WherecanIstart **Highly Voted** 1 year, 8 months ago

Selected Answer: C

Load balancer is public facing accepting all traffic coming towards the VPC (0.0.0.0/0). The web server needs to trust traffic originating from the ALB. The DB will only trust traffic originating from the Web server on port 3306 for Mysql
upvoted 6 times

TheFivePips **Most Recent** 9 months, 2 weeks ago

Selected Answer: C

Option C aligns with the least access principle and provides a clear and granular control over the communication between different components in the architecture.

Option D suggests using network ACLs, but security groups are more suitable for controlling access to individual instances based on their security group membership, which is why Option C is the more appropriate choice in this context

upvoted 3 times

TariqKipkemei 1 year, 1 month ago

Selected Answer: C

Create a security group for the web servers and allow port 443 from the load balancer. Create a security group for the MySQL servers and allow port 3306 from the web servers security group.

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

C) Create a security group for the web servers and allow port 443 from the load balancer. Create a security group for the MySQL servers and allow port 3306 from the web servers security group.

This option follows the principle of least privilege by only allowing necessary access:

Web server SG allows port 443 from load balancer SG (not open to world)

MySQL SG allows port 3306 only from web server SG

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

Create a security group for the web servers and allow port 443 from the load balancer. Create a security group for the MySQL servers and allow port 3306 from the web servers security group

upvoted 2 times

elearningtakai 1 year, 8 months ago

Selected Answer: C

Option C is the correct choice.

upvoted 2 times

fkie4 1 year, 8 months ago

Selected Answer: C

Just C. plain and simple

upvoted 2 times

aragon_saa 1 year, 8 months ago

C

<https://www.examttopics.com/discussions/amazon/view/43796-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 3 times

[Removed] 1 year, 9 months ago

Selected Answer: C

cccccc

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #386

Topic 1

An ecommerce company is running a multi-tier application on AWS. The front-end and backend tiers both run on Amazon EC2, and the database runs on Amazon RDS for MySQL. The backend tier communicates with the RDS instance. There are frequent calls to return identical datasets from the database that are causing performance slowdowns.

Which action should be taken to improve the performance of the backend?

- A. Implement Amazon SNS to store the database calls.
- B. Implement Amazon ElastiCache to cache the large datasets. **Most Voted**
- C. Implement an RDS for MySQL read replica to cache database calls.
- D. Implement Amazon Kinesis Data Firehose to stream the calls to the database.

Correct Answer: B

Community vote distribution

B (100%)

Comments

elearningtakai **Highly Voted** 1 year, 8 months ago

Selected Answer: B

the best solution is to implement Amazon ElastiCache to cache the large datasets, which will store the frequently accessed data in memory, allowing for faster retrieval times. This can help to alleviate the frequent calls to the database, reduce latency, and improve the overall performance of the backend tier.

upvoted 15 times

fruto123 **Highly Voted** 1 year, 9 months ago

Selected Answer: B

Key term is identical datasets from the database it means caching can solve this issue by cached in frequently used dataset from DB

upvoted 6 times

Manimgh **Most Recent** 1 month, 3 weeks ago

Selected Answer: B

return identical datasets = Cash

upvoted 1 times

Ucy 8 months, 1 week ago

Selected Answer: B

Answer is B This will help reduce the frequency of calls to the database and improve overall performance by serving frequently accessed data from the cache instead of fetching it from the database every time. It's not option C as it suggests implementing an RDS for MySQL read replica to cache database calls. While read replicas can offload read operations from the primary database instance and improve read scalability, they are primarily used for read scaling and high availability rather than caching. Read replicas are intended to handle read-heavy workloads by distributing read requests across multiple instances. However, they do not inherently cache data like ElastiCache does.

upvoted 4 times

Ucy 8 months, 1 week ago

Answer is B

This will help reduce the frequency of calls to the database and improve overall performance by serving frequently accessed data from the cache instead of fetching it from the database every time.

It's not option C as it suggests implementing an RDS for MySQL read replica to cache database calls. While read replicas can offload read operations from the primary database instance and improve read scalability, they are primarily used for read scaling and high availability rather than caching.

Read replicas are intended to handle read-heavy workloads by distributing read requests across multiple instances. However, they do not inherently cache data like ElastiCache does.

upvoted 2 times

Bhanu1992 8 months, 1 week ago

Keyword is identical datasets

upvoted 2 times

thewalker 10 months ago

Selected Answer: B

As per Amazon Q:

ElastiCache can be used to cache datasets from queries to RDS databases. Some key points:

While creating an ElastiCache cluster from the RDS console provides convenience, the application is still responsible for leveraging the cache.

Caching query results in ElastiCache can significantly improve performance by allowing high-volume read operations to be served from cache versus hitting the database.

This is especially useful for applications with high read throughput needs, as scaling the database can become more expensive compared to scaling the cache as needs increase. ElastiCache nodes can support up to 400,000 queries per second.

Cost savings are directly proportional to read throughput - higher throughput applications see greater savings.

upvoted 2 times

Murtadhaceit 1 year ago

Selected Answer: B

The best scenario to implement caching, identical calls to the same data sets.

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

B) Implement Amazon ElastiCache to cache the large datasets.

The key issue is repeated calls to return identical datasets from the RDS database causing performance slowdowns.

Implementing Amazon ElastiCache for Redis or Memcached would allow these repeated query results to be cached, improving backend performance by reducing load on the database.

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

B) Implement Amazon ElastiCache to cache the large datasets.

The key issue is repeated calls to return identical datasets from the RDS database causing performance slowdowns.

Implementing Amazon ElastiCache for Redis or Memcached would allow these repeated query results to be cached, improving backend performance by reducing load on the database.

upvoted 2 times

Abrar2022 1 year, 5 months ago

Selected Answer: B

Thanks Tariq for the simplified answer below:

frequent identical calls = ElastiCache

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

frequent identical calls = ElastiCache

upvoted 2 times

Mikebonsi70 1 year, 8 months ago

Tricky question, anyway.

upvoted 2 times

Mikebonsi70 1 year, 8 months ago

Yes, caching is the solution but is Elasticache compatible with RDS MySQL DB? So, what about the answer C with a DB read replica? For me it's C.

upvoted 1 times

aragon_saa 1 year, 8 months ago

B

<https://www.examtopycs.com/discussions/amazon/view/27874-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #387

Topic 1

A new employee has joined a company as a deployment engineer. The deployment engineer will be using AWS CloudFormation templates to create multiple AWS resources. A solutions architect wants the deployment engineer to perform job activities while following the principle of least privilege.

Which combination of actions should the solutions architect take to accomplish this goal? (Choose two.)

- A. Have the deployment engineer use AWS account root user credentials for performing AWS CloudFormation stack operations.
- B. Create a new IAM user for the deployment engineer and add the IAM user to a group that has the PowerUsers IAM policy attached.
- C. Create a new IAM user for the deployment engineer and add the IAM user to a group that has the AdministratorAccess IAM policy attached.
- D. Create a new IAM user for the deployment engineer and add the IAM user to a group that has an IAM policy that allows AWS CloudFormation actions only. **Most Voted**
- E. Create an IAM role for the deployment engineer to explicitly define the permissions specific to the AWS CloudFormation stack and launch stacks using that IAM role. **Most Voted**

Correct Answer: *DE*

Community vote distribution



Comments

truongtx8 **Highly Voted** 10 months, 3 weeks ago

Selected Answer: DE

The answers inside the question: CloudFormation.
A is excluded since root account is never a choice for the principle of least privilege.
D, E left are the correct ones.

upvoted 5 times

aragon_saa Highly Voted 1 year, 8 months ago

D, E

<https://www.examttopics.com/discussions/amazon/view/46428-exam-aws-certified-solutions-architect-associate-saa-c02/>
upvoted 5 times

mwwt2022 11 months, 1 week ago

thank you my friend
upvoted 2 times

awsgeek75 Most Recent 10 months, 3 weeks ago

Selected Answer: DE

ABC are just giving too much access so CD are logical choices
upvoted 3 times

TariqKipkemei 1 year, 1 month ago

Selected Answer: DE

Create a new IAM user for the deployment engineer and add the IAM user to a group that has an IAM policy that allows AWS CloudFormation actions only.
Create an IAM role for the deployment engineer to explicitly define the permissions specific to the AWS CloudFormation stack and launch stacks using that IAM role.
upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: DE

The two actions that should be taken to follow the principle of least privilege are:

D) Create a new IAM user for the deployment engineer and add the IAM user to a group that has an IAM policy that allows AWS CloudFormation actions only.

E) Create an IAM role for the deployment engineer to explicitly define the permissions specific to the AWS CloudFormation stack and launch stacks using that IAM role.

The principle of least privilege states that users should only be given the minimal permissions necessary to perform their job function.
upvoted 4 times

alexandercamachop 1 year, 6 months ago

Selected Answer: DE

Option D, creating a new IAM user and adding them to a group with an IAM policy that allows AWS CloudFormation actions only, ensures that the deployment engineer has the necessary permissions to perform AWS CloudFormation operations while limiting access to other resources and actions. This aligns with the principle of least privilege by providing the minimum required permissions for their job activities.

Option E, creating an IAM role with specific permissions for AWS CloudFormation stack operations and allowing the deployment engineer to assume that role, is another valid approach. By using an IAM role, the deployment engineer can assume the role when necessary, granting them temporary permissions to perform CloudFormation actions. This provides a level of separation and limits the permissions granted to the engineer to only the required CloudFormation operations.
upvoted 3 times

Babaaaaa 1 year, 6 months ago

Selected Answer: DE

Dddd,Eeee
upvoted 2 times

elearningtakai 1 year, 8 months ago

Selected Answer: DE

D & E are a good choices
upvoted 2 times

fruto123 1 year, 9 months ago

Selected Answer: DE

I agree DE

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #388

Topic 1

A company is deploying a two-tier web application in a VPC. The web tier is using an Amazon EC2 Auto Scaling group with public subnets that span multiple Availability Zones. The database tier consists of an Amazon RDS for MySQL DB instance in separate private subnets. The web tier requires access to the database to retrieve product information.

The web application is not working as intended. The web application reports that it cannot connect to the database. The database is confirmed to be up and running. All configurations for the network ACLs, security groups, and route tables are still in their default states.

What should a solutions architect recommend to fix the application?

- A. Add an explicit rule to the private subnet's network ACL to allow traffic from the web tier's EC2 instances.
- B. Add a route in the VPC route table to allow traffic between the web tier's EC2 instances and the database tier.
- C. Deploy the web tier's EC2 instances and the database tier's RDS instance into two separate VPCs, and configure VPC peering.
- D. Add an inbound rule to the security group of the database tier's RDS instance to allow traffic from the web tiers security group. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****TariqKipkemei** **Highly Voted** 1 year, 6 months ago**Selected Answer: D**

Security group defaults block all inbound traffic..Add an inbound rule to the security group of the database tier's RDS instance to allow traffic from the web tiers security group

upvoted 15 times

smartegnine **Highly Voted** 1 year, 5 months ago**Selected Answer: D**

Security Groups are tied on instance where as network ACL are tied to Subnet.

upvoted 5 times

ExamGuru727 Most Recent 8 months, 2 weeks ago

Selected Answer: D

For those questioning why the answer is not A:

<https://docs.aws.amazon.com/vpc/latest/userguide/vpc-network-acls.html>

Default NACLs allow all traffic, and in this question NACLs, SGs and route tables are in their default states.

upvoted 5 times

hgjdsh 8 months, 2 weeks ago

Selected Answer: A

I think the answer should be A. Since the services are in different subnets, the NACL would by default block all the incoming traffic to the subnet. Security group rule wouldn't be able to override NACL rule.

upvoted 1 times

njufi 8 months, 3 weeks ago

I selected option D as well, but I have a question regarding option A. Considering that the EC2 instances and the RDS are located in different subnets, shouldn't the network ACLs for each subnet allow traffic from one another as well? Given that the default settings for network ACLs typically block all traffic, wouldn't it be necessary to explicitly permit communication between the subnets?

upvoted 2 times

elearningtakai 1 year, 8 months ago

Selected Answer: D

By default, all inbound traffic to an RDS instance is blocked. Therefore, an inbound rule needs to be added to the security group of the RDS instance to allow traffic from the security group of the web tier's EC2 instances.

upvoted 4 times

Russs99 1 year, 8 months ago

Selected Answer: D

D is the correct answer

upvoted 2 times

aragon_saa 1 year, 8 months ago

D

<https://www.examttopics.com/discussions/amazon/view/81445-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

KAUS2 1 year, 8 months ago

Selected Answer: D

D is correct option

upvoted 2 times

[Removed] 1 year, 9 months ago

Selected Answer: D

ddddddd

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #389

Topic 1

A company has a large dataset for its online advertising business stored in an Amazon RDS for MySQL DB instance in a single Availability Zone. The company wants business reporting queries to run without impacting the write operations to the production DB instance.

Which solution meets these requirements?

- A. Deploy RDS read replicas to process the business reporting queries. **Most Voted**
- B. Scale out the DB instance horizontally by placing it behind an Elastic Load Balancer.
- C. Scale up the DB instance to a larger instance type to handle write operations and queries.
- D. Deploy the DB instance in multiple Availability Zones to process the business reporting queries.

Correct Answer: A

Community vote distribution

A (100%)

Comments

mwwt2022 **Highly Voted** 11 months, 1 week ago

Selected Answer: A

reporting queries to run without impacting the write operations -> read replicas
upvoted 5 times

Guru4Cloud **Most Recent** 1 year, 3 months ago

Selected Answer: A

A) Deploy RDS read replicas to process the business reporting queries.

The key points are:

RDS read replicas allow read-only copies of the production DB instance to be created
Queries to the read replica don't affect the source DB instance performance
This isolates reporting queries from production traffic and write operations
So using RDS read replicas is the best way to meet the requirements of running reporting queries without impacting production write operations.
upvoted 4 times

james2033 1 year, 4 months ago

Selected Answer: A

"single AZ", "large dataset", "Amazon RDS for MySQL database". Want "business report queries". --> Solution "Read replicas", choose A.

upvoted 2 times

antropaws 1 year, 6 months ago

Selected Answer: A

No doubt A.

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Load balance read operations = read replicas

upvoted 2 times

TariqKipkemei 1 year, 1 month ago

reports=read replica

upvoted 2 times

KAUS2 1 year, 8 months ago

Selected Answer: A

Option "A" is the right answer . Read replica use cases - You have a production database that is taking on normal load & You want to run a reporting application to run some analytics

- You create a Read Replica to run the new workload there
- The production application is unaffected
- Read replicas are used for SELECT (=read) only kind of statements (not INSERT, UPDATE, DELETE)

upvoted 3 times

[Removed] 1 year, 9 months ago

Selected Answer: A

aaaaaaaaaa

upvoted 2 times

cegama543 1 year, 9 months ago

Selected Answer: A

option A is the best solution for ensuring that business reporting queries can run without impacting write operations to the production DB instance.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #390

Topic 1

A company hosts a three-tier ecommerce application on a fleet of Amazon EC2 instances. The instances run in an Auto Scaling group behind an Application Load Balancer (ALB). All ecommerce data is stored in an Amazon RDS for MariaDB Multi-AZ DB instance.

The company wants to optimize customer session management during transactions. The application must store session data durably.

Which solutions will meet these requirements? (Choose two.)

- A. Turn on the sticky sessions feature (session affinity) on the ALB. **Most Voted**
- B. Use an Amazon DynamoDB table to store customer session information.
- C. Deploy an Amazon Cognito user pool to manage user session information.
- D. Deploy an Amazon ElastiCache for Redis cluster to store customer session information. **Most Voted**
- E. Use AWS Systems Manager Application Manager in the application to manage user session information.

Correct Answer: AD

Community vote distribution



Comments

fruto123 **Highly Voted** 2 years, 2 months ago

Selected Answer: AD

It is A and D. Proof is in link below.

<https://aws.amazon.com/caching/session-management/>
upvoted 28 times

babayomi 7 months, 2 weeks ago

Thanks,
First consider the durability requirement of a session. It's never meant to be kept after the session terminates, tha5 said. You wouldn't need a permanent database for a session. You would only require a cache, so that it can expire after a while when the session has already completed. That is why D is correct. A is required in the first place in order to initialize the use of the

session has already completed. That is why D is correct. A is required in the first place in order to initialize the use of the cache.

upvoted 2 times

pentium75 1 year, 5 months ago

This doesn't say anything about durability

upvoted 4 times

babayomi 7 months, 2 weeks ago

First consider the durability requirement of a session. It's never meant to be kept after the session terminates, tha5 said. You wouldn't need a pamanent database for a session. You would only require a cache, so that it can expire after a while when the session has already completed. That is why D is correct. A is required in the first place in order to initialize the use of the cache.

upvoted 2 times

Marco_St Highly Voted 1 year, 5 months ago

Selected Answer: BD

I did not get why A is most voted? The question did not mention anything about fixed routing target so the ALB should route traffic randomly to each server. Then we just need to provide cache session management to avoid session lost issue instead of using sticky session.

upvoted 14 times

babayomi 7 months, 2 weeks ago

AD

First consider the durability requirement of a session. It's never meant to be kept after the session terminates, tha5 said. You wouldn't need a pamanent database for a session. You would only require a cache, so that it can expire after a while when the session has already completed. That is why D is correct. A is required in the first place in order to initialize the use of the cache.

upvoted 1 times

chest_jd Most Recent 1 month, 4 weeks ago

Selected Answer: BD

First of all - application MUST STORE session DATA. We should get some storage on backend side. Just Band D can do it.

upvoted 1 times

bigift 3 months, 1 week ago

Selected Answer: BD

Avoid sticky sessions:

While using sticky sessions on an Application Load Balancer (ALB) might seem like a solution, it can hinder scalability as it ties a user to a specific instance, which could become a bottleneck if traffic spikes.

Dynamodb and Elastic cache are best i think

upvoted 1 times

zdi561 4 months ago

Selected Answer: BD

A is not right, the question is to where to save user session, the sticky session is used for ALB to send the request of the same client to the same target.

upvoted 1 times

FlyingHawk 6 months, 1 week ago

Selected Answer: BD

People select AD because they think it is combine A and D to make a final solution. if the expectation is that each selected option should independently meet the requirements, then B and D. I think "Which solutions will meet these requirements?" should mean two independent solutions.

upvoted 2 times

Arad 6 months, 2 weeks ago

Selected Answer: BD

session management during transaction ---> DynamoDB and ElastiCache

unvoted 1 times

upvoted 1 times

Sergantus 7 months, 1 week ago

Selected Answer: BD

We are not looking here for a combination of options, just two independent solutions.

Because of the "durability" requirement -

B - DynamoDB is durable

D - ElastiCache for Redis although mostly used as in-memory cache, ElastiCache for Redis provides options for data persistence - RDB (Redis Database Backup) and AOF (Append-Only File). RDB periodically saves a snapshot of the dataset to disk, while AOF logs every write operation and replays it in case of a restart. ElastiCache for Redis also allows configuring replication groups with Multi-AZ and read replicas, providing additional redundancy and failover options.

upvoted 3 times

FlyingHawk 6 months, 1 week ago

I agree with your comments

upvoted 1 times

babayomi 7 months, 2 weeks ago

AD

First consider the durability requirement of a session. It's never meant to be kept after the session terminates, tha5 said. You wouldn't need a pamanent database for a session. You would only require a cache, so that it can expire after a while when the session has already completed. That is why D is correct. A is required in the first place in order to initialize the use of the cache.

upvoted 1 times

Uzbekistan 1 year, 3 months ago

Selected Answer: BD

Option A suggests using sticky sessions (session affinity) on the Application Load Balancer (ALB). While sticky sessions can help route requests from the same client to the same backend server, it doesn't directly address the requirement for durable storage of session data. Sticky sessions are typically used to maintain session state at the load balancer level, but they do not provide data durability in case of server failures or restarts.

Option A - is not correct !!!

So answer is option B and D !!!

upvoted 3 times

babayomi 7 months, 2 weeks ago

First consider the durability requirement of a session. It's never meant to be kept after the session terminates, tha5 said. You wouldn't need a pamanent database for a session. You would only require a cache, so that it can expire after a while when the session has already completed. That is why D is correct. A is required in the first place in order to initialize the use of the cache.

upvoted 1 times

jjcode 1 year, 3 months ago

why does it matter to store user sessions durably? they EXPIRE, why would a company care about storing user sessions, thats not something thats done in the real world, those things are usually data dumped, or overwritten with new session tokens LOL, this whole question is &^%&*^\$#@%^

upvoted 4 times

tuso 1 year, 4 months ago

I think the question is intended to mean "Combination of services", as some answers say "to store" or "to manage". So i am going for A+B, as sticky sessions are intended to manage the sessions and DynamoDB to store durably.

upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: AB

Going for AB. Sticky Sessions to "optimize customer session management during transactions" and DynamoDB to "store session data durably".

D, ElastiCache does NOT allow "durable" storage. Just because there's an article that contains both words "ElastiCache" and "durable" does not prove the contrary.

C and E, Cognito and Systems Manager, have nothing to do with the issue.

upvoted 5 times

dkw2342 1 year, 3 months ago

I agree that ElastiCache for Redis is not a durable KV store.

But what about the phrasing?

"Which solutions will meet these requirements? (Choose two.)" Solutions (plural) implies two ways to *independently* fulfill the requirements. If you're supposed to select a combination of options, it's usually phrased like this: "Which combination of solutions ..."

upvoted 4 times

avdxeqtr 1 year, 5 months ago

<https://aws.amazon.com/blogs/developer/elasticache-as-an-asp-net-session-store/>

Amazon ElastiCache for Redis is highly suited as a session store to manage session information such as user authentication tokens, session state, and more. Simply use ElastiCache for Redis as a fast key-value store with appropriate TTL on session keys to manage your session information. Session management is commonly required for online applications, including games, e-commerce websites, and social media platforms.

upvoted 2 times

avdxeqtr 1 year, 5 months ago

Correct link: <https://aws.amazon.com/elasticache/redis/>

upvoted 3 times

m_y_s 1 year, 6 months ago

Selected Answer: BD

I don't understand what Sticky Session has to do with session storage. For the intent of the problem, I think DynamoDB and Redis are appropriate.

upvoted 6 times

pentium75 1 year, 5 months ago

"Session storage" is not the only requirement here. It is about 'optimizing customer session management during transactions', obviously it makes sense to host customer sessions on same node to easy the session management.

upvoted 4 times

Sergantus 7 months, 1 week ago

Up until that node fails

upvoted 2 times

daniel1 1 year, 7 months ago

Selected Answer: BD

Chatgpt4 says B and D

Option A (Sticky sessions) is more for ensuring that a client's requests are sent to the same target once a session is established, but it doesn't provide a mechanism for durable session data storage across multiple instances. Option C (Amazon Cognito) is more for user identity management rather than session data storage during transactions. Option E (AWS Systems Manager Application Manager) is not a suitable or standard choice for session management in applications.

upvoted 4 times

pentium75 1 year, 5 months ago

Answers starting with "ChatGPT says ..." are usually wrong.

In that case, B and D solve the same part of the requirement (storing session data), just B is durable (as required) while D is not durable (thus failing to meet the requirement). We still need to 'optimize customer session management'.

upvoted 4 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: AD

Well, this documentation says it all. Option A is obvious, and D ElastiCache for Redis, can even support replication in case of node failure/session data loss

node failure/session data loss.

<https://aws.amazon.com/caching/session-management/>

upvoted 4 times

pentium75 1 year, 5 months ago

ElastiCache can be HA and supports replication, but it remains a cache, which is by definition not durable.

upvoted 3 times

Sergantus 7 months, 1 week ago

You can configure Redis with persistence options

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: AD

It is A and D. Proof is in link below.

<https://aws.amazon.com/caching/session-management/>

upvoted 2 times

pentium75 1 year, 5 months ago

That does not say anything about durability.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #391

Topic 1

A company needs a backup strategy for its three-tier stateless web application. The web application runs on Amazon EC2 instances in an Auto Scaling group with a dynamic scaling policy that is configured to respond to scaling events. The database tier runs on Amazon RDS for PostgreSQL. The web application does not require temporary local storage on the EC2 instances. The company's recovery point objective (RPO) is 2 hours.


The backup strategy must maximize scalability and optimize resource utilization for this environment.

Which solution will meet these requirements?

- A. Take snapshots of Amazon Elastic Block Store (Amazon EBS) volumes of the EC2 instances and database every 2 hours to meet the RPO.
- B. Configure a snapshot lifecycle policy to take Amazon Elastic Block Store (Amazon EBS) snapshots. Enable automated backups in Amazon RDS to meet the RPO.
- C. Retain the latest Amazon Machine Images (AMIs) of the web and application tiers. Enable automated backups in Amazon RDS and use point-in-time recovery to meet the RPO. **Most Voted**
- D. Take snapshots of Amazon Elastic Block Store (Amazon EBS) volumes of the EC2 instances every 2 hours. Enable automated backups in Amazon RDS and use point-in-time recovery to meet the RPO.

Correct Answer: C

Community vote distribution

**Comments**

elearningtakai **Highly Voted** 2 years, 2 months ago

Selected Answer: C

that if there is no temporary local storage on the EC2 instances, then snapshots of EBS volumes are not necessary. Therefore, if your application does not require temporary storage on EC2 instances, using AMIs to back up the web and application tiers is sufficient to restore the system after a failure.

Snapshots of EBS volumes would be necessary if you want to back up the entire EC2 instance, including any applications and temporary data stored on the EBS volumes attached to the instances. When you take a snapshot of an EBS volume, it backs up the entire contents of that volume. This ensures that you can restore the entire EC2 instance to a specific point in time more quickly. However, if there is no temporary data stored on the EBS volumes, then snapshots of EBS volumes are not necessary.

quickly. However, if there is no temporary data stored on the EBS volumes, then snapshots of EBS volumes are not necessary.
upvoted 36 times

MssP 2 years, 2 months ago

I think "temporal local storage" refers to "instance store", no instance store is required. EBS is durable storage, not temporal.
upvoted 4 times

MssP 2 years, 2 months ago

Look at the first paragraph. <https://repost.aws/knowledge-center/instance-store-vs-ebs>
upvoted 2 times

MatAlves 9 months ago

Considering it's a "stateless web application", that would still be no reason to back up the EBS volumes.
upvoted 2 times

CloudForFun Highly Voted 2 years, 2 months ago

Selected Answer: C

The web application does not require temporary local storage on the EC2 instances => No EBS snapshot is required, retaining the latest AMI is enough.
upvoted 15 times

Mikado211 Most Recent 1 year, 5 months ago

Selected Answer: C

The web application does not require temporary local storage on the EC2 instances so we do not care about ECS. We only need two things here, the image of the instance (AMI) and a database backup.

C
upvoted 4 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: C

"The web application does not require temporary local storage on the EC2 instances" rules out any option to back up the EC2 EBS volumes.
upvoted 3 times

darekw 1 year, 10 months ago

Question says: ...stateless web application.. that means application doesn't store any data, so no EBS required
upvoted 2 times

kruasan 2 years, 1 month ago

Selected Answer: C

Since the application has no local data on instances, AMIs alone can meet the RPO by restoring instances from the most recent AMI backup. When combined with automated RDS backups for the database, this provides a complete backup solution for this environment.

The other options involving EBS snapshots would be unnecessary given the stateless nature of the instances. AMIs provide all the backup needed for the app tier.

This uses native, automated AWS backup features that require minimal ongoing management:

- AMI automated backups provide point-in-time recovery for the stateless app tier.
- RDS automated backups provide point-in-time recovery for the database.

upvoted 4 times

neosis91 2 years, 1 month ago

Selected Answer: B

BBBBBBBBBB
upvoted 1 times

pentium75 1 year, 5 months ago

Why back up EBS volumes of the autoscaled instances?

upvoted 2 times

Rob1L 2 years, 2 months ago

Selected Answer: D

I vote for D

upvoted 1 times

pentium75 1 year, 5 months ago

Why back up EBS volumes of the autoscaled instances?

upvoted 2 times

CapJackSparrow 2 years, 2 months ago

Selected Answer: C

makes more sense.

upvoted 3 times

nileshlg 2 years, 2 months ago

Selected Answer: C

Answer is C. Keyword to notice "Stateless"

upvoted 3 times

cra2yk 2 years, 2 months ago

Selected Answer: C

why B? I mean "stateless" and "does not require temporary local storage" have indicate that we don't need to take snapshot for ec2 volume.

upvoted 4 times

ktulu2602 2 years, 2 months ago

Selected Answer: B

Option B is the most appropriate solution for the given requirements.

With this solution, a snapshot lifecycle policy can be created to take Amazon Elastic Block Store (Amazon EBS) snapshots periodically, which will ensure that EC2 instances can be restored in the event of an outage. Additionally, automated backups can be enabled in Amazon RDS for PostgreSQL to take frequent backups of the database tier. This will help to minimize the RPO to 2 hours.

Taking snapshots of Amazon EBS volumes of the EC2 instances and database every 2 hours (Option A) may not be cost-effective and efficient, as this approach would require taking regular backups of all the instances and volumes, regardless of whether any changes have occurred or not. Retaining the latest Amazon Machine Images (AMIs) of the web and application tiers (Option C) would provide only an image backup and not a data backup, which is required for the database tier. Taking snapshots of Amazon EBS volumes of the EC2 instances every 2 hours and enabling automated backups in Amazon RDS and using point-in-time recovery (Option D) would result in higher costs and may not be necessary to meet the RPO requirement of 2 hours.

upvoted 4 times

pentium75 1 year, 5 months ago

Why back up EBS volumes of the autoscaled instances?

"Retaining the latest Amazon Machine Images (AMIs) of the web and application tiers (Option C) would provide only an image backup and not a data backup, which is required for the database tier." False because option C also includes "automated backups in Amazon RDS".

upvoted 2 times

cegama543 2 years, 3 months ago

Selected Answer: B

B. Configure a snapshot lifecycle policy to take Amazon Elastic Block Store (Amazon EBS) snapshots. Enable automated backups in Amazon RDS to meet the RPO.

The best solution is to configure a snapshot lifecycle policy to take Amazon Elastic Block Store (Amazon EBS) snapshots, and

enable automated backups in Amazon RDS to meet the RPO. An RPO of 2 hours means that the company needs to ensure that the backup is taken every 2 hours to minimize data loss in case of a disaster. Using a snapshot lifecycle policy to take Amazon EBS snapshots will ensure that the web and application tier can be restored quickly and efficiently in case of a disaster. Additionally, enabling automated backups in Amazon RDS will ensure that the database tier can be restored quickly and efficiently in case of a disaster. This solution maximizes scalability and optimizes resource utilization because it uses automated backup solutions built into AWS.

upvoted 3 times

pentium75 1 year, 5 months ago

No need to back up the EBS volumes of autoscaled instances.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #392

Topic 1

A company wants to deploy a new public web application on AWS. The application includes a web server tier that uses Amazon EC2 instances. The application also includes a database tier that uses an Amazon RDS for MySQL DB instance.

The application must be secure and accessible for global customers that have dynamic IP addresses.

How should a solutions architect configure the security groups to meet these requirements?


A. Configure the security group for the web servers to allow inbound traffic on port 443 from 0.0.0.0/0. Configure the security group for the DB instance to allow inbound traffic on port 3306 from the security group of the web servers.

Most Voted

B. Configure the security group for the web servers to allow inbound traffic on port 443 from the IP addresses of the customers. Configure the security group for the DB instance to allow inbound traffic on port 3306 from the security group of the web servers.

C. Configure the security group for the web servers to allow inbound traffic on port 443 from the IP addresses of the customers. Configure the security group for the DB instance to allow inbound traffic on port 3306 from the IP addresses of the customers.

D. Configure the security group for the web servers to allow inbound traffic on port 443 from 0.0.0.0/0. Configure the security group for the DB instance to allow inbound traffic on port 3306 from 0.0.0.0/0.

Correct Answer: A*Community vote distribution***Comments****Bhanu1992** **Highly Voted** 8 months, 1 week ago

The keyword is dynamic IPs from the customer, then B, C out, D out due to 0.0.0.0/0
upvoted 6 times

awsgeek75 **Highly Voted** 10 months, 3 weeks ago**Selected Answer: A**

"The application must be secure and accessible for global customers that have dynamic IP addresses." This just means "anyone"

so BC are wrong as you cannot know in advance about the dynamic IP addresses. D is just opening the DB to the internet.

A is most secure as web is open to internet and db is open to web only.

upvoted 6 times

bora4motion Most Recent 1 month ago

Selected Answer: A

a is 100% correct.

upvoted 1 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

It allows HTTPS access from any public IP address, meeting the requirement for global customer access.

HTTPS provides encryption for secure communication.

And for the database security group, only allowing inbound port 3306 from the web server security group properly restricts access to only the resources that need it.

upvoted 4 times

jayce5 1 year, 6 months ago

Selected Answer: A

Should be A since the customer IPs are dynamically.

upvoted 2 times

antropaws 1 year, 6 months ago

Selected Answer: A

A no doubt.

upvoted 3 times

omoakin 1 year, 6 months ago

BBBBBBBBBBBBBBBBBBBBBB

from customers IPs

upvoted 1 times

MostafaWardany 1 year, 5 months ago

Correct answer A, customer dynamic IPs ==>> 443 from 0.0.0.0/0

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

dynamic source ips = allow all traffic - Configure the security group for the web servers to allow inbound traffic on port 443 from 0.0.0.0/0. Configure the security group for the DB instance to allow inbound traffic on port 3306 from the security group of the web servers.

upvoted 2 times

elearningtakai 1 year, 8 months ago

Selected Answer: A

If the customers have dynamic IP addresses, option A would be the most appropriate solution for allowing global access while maintaining security.

upvoted 4 times

Kenzo 1 year, 8 months ago

Correct answer is A.

B and C are out.

D is out because it is accepting traffic from every where instead of from webservers only

upvoted 4 times

Grace83 1 year, 8 months ago

A is correct

upvoted 3 times

Whericanlstart 1 year, 8 months ago

Selected Answer: B

Keyword dynamic ...A is the right answer. If the IP were static and specific, B would be the right answer

upvoted 4 times

pentium75 11 months, 2 weeks ago

Then why voted B?

upvoted 2 times

boxu03 1 year, 8 months ago

Selected Answer: A

aaaaaaa

upvoted 1 times

kprakashbehera 1 year, 9 months ago

Selected Answer: A

Ans - A

upvoted 1 times

[Removed] 1 year, 9 months ago

Selected Answer: A

aaaaaa

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #393

Topic 1

A payment processing company records all voice communication with its customers and stores the audio files in an Amazon S3 bucket. The company needs to capture the text from the audio files. The company must remove from the text any personally identifiable information (PII) that belongs to customers.

What should a solutions architect do to meet these requirements?

- A. Process the audio files by using Amazon Kinesis Video Streams. Use an AWS Lambda function to scan for known PII patterns.
- B. When an audio file is uploaded to the S3 bucket, invoke an AWS Lambda function to start an Amazon Textract task to analyze the call recordings.
- C. Configure an Amazon Transcribe transcription job with PII redaction turned on. When an audio file is uploaded to the S3 bucket, invoke an AWS Lambda function to start the transcription job. Store the output in a separate S3 bucket. **Most Voted**
- D. Create an Amazon Connect contact flow that ingests the audio files with transcription turned on. Embed an AWS Lambda function to scan for known PII patterns. Use Amazon EventBridge to start the contact flow when an audio file is uploaded to the S3 bucket.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**TariqKipkemei** **Highly Voted** 7 months, 2 weeks ago**Selected Answer: C**

speech to text = Amazon Transcribe
upvoted 5 times

Guru4Cloud **Highly Voted** 9 months, 1 week ago**Selected Answer: C**

Amazon Transcribe is a service provided by Amazon Web Services (AWS) that converts speech to text using automatic speech recognition (ASR) technology
upvoted 5 times

james2033 Most Recent 10 months, 3 weeks ago

Selected Answer: C

AWS Transcribe <https://aws.amazon.com/transcribe/> . Redacting or identifying (Personally identifiable instance) PII in real-time stream <https://docs.aws.amazon.com/transcribe/latest/dg/pii-redaction-stream.html> .

upvoted 2 times

SimiTik 1 year, 1 month ago

C

Amazon Transcribe is a service provided by Amazon Web Services (AWS) that converts speech to text using automatic speech recognition (ASR) technology. gtp

upvoted 4 times

elearningtakai 1 year, 2 months ago

Selected Answer: C

Option C is the most suitable solution as it suggests using Amazon Transcribe with PII redaction turned on. When an audio file is uploaded to the S3 bucket, an AWS Lambda function can be used to start the transcription job. The output can be stored in a separate S3 bucket to ensure that the PII redaction is applied to the transcript. Amazon Transcribe can redact PII such as credit card numbers, social security numbers, and phone numbers.

upvoted 4 times

WheracanIstart 1 year, 2 months ago

Selected Answer: C

C for sure.....

upvoted 2 times

WheracanIstart 1 year, 2 months ago

C for sure

upvoted 2 times

boxu03 1 year, 2 months ago

Selected Answer: C

cccccccc

upvoted 2 times

Ruhi02 1 year, 2 months ago

answer c

upvoted 2 times

KAUS2 1 year, 2 months ago

Selected Answer: C

Option C is correct..

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #394

Topic 1

A company is running a multi-tier ecommerce web application in the AWS Cloud. The application runs on Amazon EC2 instances with an Amazon RDS for MySQL Multi-AZ DB instance. Amazon RDS is configured with the latest generation DB instance with 2,000 GB of storage in a General Purpose SSD (gp3) Amazon Elastic Block Store (Amazon EBS) volume. The database performance affects the application during periods of high demand.

A database administrator analyzes the logs in Amazon CloudWatch Logs and discovers that the application performance always degrades when the number of read and write IOPS is higher than 20,000.

What should a solutions architect do to improve the application performance?

- A. Replace the volume with a magnetic volume.
- B. Increase the number of IOPS on the gp3 volume.
- C. Replace the volume with a Provisioned IOPS SSD (io2) volume. **Most Voted**
- D. Replace the 2,000 GB gp3 volume with two 1,000 GB gp3 volumes.

Correct Answer: C

Community vote distribution



Comments

Bezha **Highly Voted** 2 years, 2 months ago

Selected Answer: D

A - Magnetic Max IOPS 200 - Wrong
 B - gp3 Max IOPS 16000 per volume - Wrong
 C - RDS not supported io2 - Wrong
 D - Correct; 2 gp3 volume with 16 000 each $2 \times 16000 = 32\ 000$ IOPS
 upvoted 37 times

FlyingHawk 6 months, 1 week ago

gp3 Max IOPS 16000 per volume is for SQL server, for MySQL server it is 64000 per volume, so B should work.
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html
 upvoted 1 times

dkw2342 1 year, 3 months ago

I really wonder how this answer can be the top answer. How would it even be possible to provision multiple gp3 volumes for RDS? RDS manages the storage, we have no influence on the number of volumes.

Striping is something that RDS does automatically depending on storage class and volume size: "When you select General Purpose SSD or Provisioned IOPS SSD, depending on the engine selected and the amount of storage requested, Amazon RDS automatically stripes across multiple volumes to enhance performance (...)"

For MariaDB with 400 to 64,000 GiB of gp3 storage, RDS automatically provisions 4 volumes. This gives us 12,000 IOPS *baseline* and can be increased up to 64,000 *provisioned* IOPS.

RDS does not support io2.

Therefore: Option B
upvoted 6 times

dkw2342 1 year, 3 months ago

PS: https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html
upvoted 2 times

zits88 1 year, 2 months ago

It must be that io2 was originally not supported by RDS, because I see this untruth reposted everywhere. It totally is.
upvoted 4 times

joechen2023 1 year, 11 months ago

<https://repost.aws/knowledge-center/ebs-volume-type-differences>
RDS does support io2
upvoted 3 times

wRhIH 1 year, 11 months ago

that Link is to EBS instead of RDS
upvoted 6 times

baba365 1 year, 8 months ago

'the application performance always degrades when the number of read and write IOPS is higher than 20,000' ... question didn't say read and write IOPS can't be higher than 32,000. Answer: C if it's based on performance and not cost related.

'Amazon RDS provides three storage types: General Purpose SSD (also known as gp2 and gp3), Provisioned IOPS SSD (also known as io1), and magnetic (also known as standard). They differ in performance characteristics and price, which means that you can tailor your storage performance and cost to the needs of your database workload.'

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html
upvoted 1 times

Michal_L_95 Highly Voted 2 years, 2 months ago

Selected Answer: B

It can not be option C as RDS does not support io2 storage type (only io1).

Here is a link to the RDS storage documentation:

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html

Also it is not the best option to take Magnetic storage as it supports max 1000 IOPS.

I vote for option B as gp3 storage type supports up to 64 000 IOPS where question mentioned with problem at level of 20 000.

upvoted 15 times

joechen2023 1 year, 11 months ago

check the link below <https://repost.aws/knowledge-center/ebs-volume-type-differences>
it states:

General Purpose SSD volumes are good for a wide variety of transactional workloads that require less than the following:

16,000 IOPS
1,000 MiB/s of throughput
160-TiB volume size

upvoted 2 times

upvoted 2 times

GalileoEC2 2 years, 2 months ago

is this true? Amazon RDS (Relational Database Service) supports the Provisioned IOPS SSD (io2) storage type for its database instances. The io2 storage type is designed to deliver predictable performance for critical and highly demanding database workloads. It provides higher durability, higher IOPS, and lower latency compared to other Amazon EBS (Elastic Block Store) storage types. RDS offers the option to choose between the General Purpose SSD (gp3) and Provisioned IOPS SSD (io2) storage types for database instances.

upvoted 3 times

1rob 1 year, 4 months ago

Please add a reference where it states that io2 is supported by RDS.

upvoted 1 times

zits88 1 year, 2 months ago

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html - Right there in the first paragraph: "Amazon RDS provides three storage types: General Purpose SSD (also known as gp2 and gp3), Provisioned IOPS SSD (also known as io1 and io2 Block Express), and magnetic (also known as standard). They differ in performance characteristics and price, which means that you can tailor your storage performance and cost to the needs of your database workload. You can create Db2, MySQL, MariaDB, Oracle, and PostgreSQL RDS DB instances with up to 64 tebibytes (TiB) of storage. "

upvoted 7 times

zdi561 Most Recent 4 months ago

Selected Answer: C

RDS now support io2

upvoted 2 times

EllenLiu 5 months, 3 weeks ago

Selected Answer: C

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html#Concepts.Storage.GeneralSSD

upvoted 1 times

FlyingHawk 6 months, 1 week ago

Selected Answer: B

I selected C initially, but after read the doc https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html, noticed maximum IOPS for gp3 volume is 64K (16K for SQL server), For an RDS for MySQL workload, increasing the provisioned IOPS on the gp3 volume to exceed the application's peak demand (e.g., 30,000 IOPS) is a cost-effective and scalable solution.

upvoted 1 times

LeonSauveterre 6 months, 1 week ago

Selected Answer: C

How on earth would the answer be D? Splitting the gp3 volume into two does not solve the IOPS limitation because each volume would still be subject to the gp3 performance cap. The core of the crux is still there!

C - io2 volumes are exactly designed for workloads requiring consistent, high-performance storage.

upvoted 1 times

Mish 6 months, 4 weeks ago

Selected Answer: C

Provisioned IOPS SSDs (io2) are specifically designed to deliver sustained high performance and low latency (RDS is supported in IO2). They can handle more than 20,000 IOPS.

upvoted 2 times

studydue 7 months ago

Updated 2024

Answer: C

upvoted 4 times

ashishs174 8 months, 2 weeks ago

Answer is C, io2 volumes are supported

<https://aws.amazon.com/blogs/aws/amazon-rds-now-supports-io2-block-express-volumes-for-mission-critical-database-workloads/>

upvoted 2 times

MatAlves 9 months ago

Nice to see that everyone just picked a different answer...

upvoted 3 times

ChymKuBoy 9 months, 3 weeks ago

Selected Answer: B

B for sure

upvoted 1 times

example_ 10 months, 1 week ago

Selected Answer: C

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html

upvoted 5 times

FrozenCarrot 11 months ago

Selected Answer: C

Now EBS support io2.

upvoted 6 times

theamachine 11 months, 3 weeks ago

Selected Answer: C

Provisioned IOPS SSDs (io2) are specifically designed to deliver sustained high performance and low latency (RDS is supported in IO2). They can handle more than 20,000 IOPS.

upvoted 7 times

FlyingHawk 6 months, 1 week ago

C is more expensive and overkill.

upvoted 1 times

Lin878 11 months, 4 weeks ago

Selected Answer: C

It should be "C" right, now.

<https://aws.amazon.com/blogs/aws/amazon-rds-now-supports-io2-block-express-volumes-for-mission-critical-database-workloads/>

upvoted 5 times

[Removed] 12 months ago

C is the correct one

EBS Volume Types Use cases

Provisioned IOPS (PIOPS) SSD

- Critical business applications with sustained IOPS performance
- Or applications that need more than 16,000 IOPS
- Great for databases workloads (sensitive to storage perf and consistency)
- io1/io2 (4 GiB - 16 TiB):
 - Max PIOPS: 64,000 for Nitro EC2 instances & 32,000 for other
 - Can increase PIOPS independently from storage size
- io2 have more durability and more IOPS per GiB (at the same price as io1)
- io2 Block Express (4 GiB - 64 TiB):
 - Sub-millisecond latency
 - Max PIOPS: 256,000 with an IOPS:GiB ratio of 1,000:1

upvoted 3 times

Scheldon 1 year ago

Selected Answer: C

Per the newest info it should be C right now

for the new database instance as follows:

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #395

Topic 1

An IAM user made several configuration changes to AWS resources in their company's account during a production deployment last week. A solutions architect learned that a couple of security group rules are not configured as desired. The solutions architect wants to confirm which IAM user was responsible for making changes.

Which service should the solutions architect use to find the desired information?

- A. Amazon GuardDuty
- B. Amazon Inspector
- C. AWS CloudTrail **Most Voted**
- D. AWS Config

Correct Answer: C*Community vote distribution*

C (100%)

Comments**cegama543** **Highly Voted** 1 year, 9 months ago**Selected Answer: C**

C. AWS CloudTrail

The best option is to use AWS CloudTrail to find the desired information. AWS CloudTrail is a service that enables governance, compliance, operational auditing, and risk auditing of AWS account activities. CloudTrail can be used to log all changes made to resources in an AWS account, including changes made by IAM users, EC2 instances, AWS management console, and other AWS services. By using CloudTrail, the solutions architect can identify the IAM user who made the configuration changes to the security group rules.

upvoted 14 times

Manimgh **Most Recent** 1 month, 3 weeks ago**Selected Answer: C**

which IAM user was responsible = CloudTrail

upvoted 1 times

BatVanyo 7 months, 3 weeks ago**Selected Answer: C**

I was initially a bit confused on what Config and CloudTrail actually do, as both can be used to track configuration changes. However, this explanation is probably the best one I have come across so far:
"Config reports on what has changed, whereas CloudTrail reports on who made the change, when, and from which location"

Since the question is which IAM user was responsible for making the changes, the answer is CloudTrail.
upvoted 4 times

d401c0d 8 months, 1 week ago

Selected Answer: C

CloudTrail = which user made which api calls. This is used for audit purpose.
upvoted 3 times

sheq 12 months ago

This question is the same with the question 388, isn't it?
upvoted 1 times

kambarami 1 year, 2 months ago

This is how you know not to trust the moderators with their answers.
upvoted 1 times

Wayne23Fang 1 year, 3 months ago

There is an article "How to use AWS Config and CloudTrail to find who made changes to a resource" in aws blog. Given CloudTrail provided AWS config original info, it seems for this particular one, C is better than AWS config.
upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

AWS CloudTrail is the correct service to use here to identify which user was responsible for the security group configuration changes
upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

AWS CloudTrail
upvoted 2 times

Bezha 1 year, 8 months ago

Selected Answer: C

AWS CloudTrail
upvoted 2 times

[Removed] 1 year, 8 months ago

Selected Answer: C

C. AWS CloudTrail
upvoted 2 times

kprakashbehera 1 year, 9 months ago

Selected Answer: C

CloudTrail logs will tell who did that
upvoted 3 times

KAUS2 1 year, 9 months ago

Selected Answer: C

Option "C" AWS CloudTrail is correct.
upvoted 3 times

Nithin1119 1 year, 9 months ago

cccccc

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #396

Topic 1

A company has implemented a self-managed DNS service on AWS. The solution consists of the following:

- Amazon EC2 instances in different AWS Regions
- Endpoints of a standard accelerator in AWS Global Accelerator

The company wants to protect the solution against DDoS attacks.

What should a solutions architect do to meet this requirement?

- A. Subscribe to AWS Shield Advanced. Add the accelerator as a resource to protect. **Most Voted**
- B. Subscribe to AWS Shield Advanced. Add the EC2 instances as resources to protect.
- C. Create an AWS WAF web ACL that includes a rate-based rule. Associate the web ACL with the accelerator.
- D. Create an AWS WAF web ACL that includes a rate-based rule. Associate the web ACL with the EC2 instances.

Correct Answer: A

Community vote distribution



Comments

WherecanIstart **Highly Voted** 1 year, 8 months ago

Selected Answer: A

DDoS attacks = AWS Shield Advance
 Shield Advance protects Global Accelerator, NLB, ALB, etc
 upvoted 13 times

Abrar2022 **Highly Voted** 1 year, 5 months ago

Selected Answer: A

DDoS attacks = AWS Shield Advance
 resource as Global Acc
 upvoted 5 times

pentium75 **Most Recent** 11 months, 2 weeks ago

Selected Answer: A

Global Accelerator is what is exposed to the Internet = where DDoS attacks could land = what must be protected by Shield Advanced

upvoted 5 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

So, the correct option is:

B. Subscribe to AWS Shield Advanced. Add the EC2 instances as resources to protect.

Here's why this option is the most appropriate:

A. While you can add the accelerator as a resource to protect with AWS Shield Advanced, it's generally more effective to protect the individual resources (in this case, the EC2 instances) because AWS Shield Advanced will automatically protect resources associated with Global Accelerator

upvoted 1 times

awsgeek75 10 months, 3 weeks ago

Which EC2 instance? Global Accelerator works by providing anycast IP addresses for the underlying resource (our EC2 in this case) so every end user trying to reach the EC2 server HAS to go through the Global Accelerator which is why the Global Accelerator needs to be protected and not the EC2.

upvoted 5 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

DDoS attacks = AWS Shield Advanced

upvoted 3 times

nileshlg 1 year, 8 months ago

Selected Answer: A

Answer is A

<https://docs.aws.amazon.com/waf/latest/developerguide/ddos-event-mitigation-logic-gax.html>

upvoted 2 times

ktulu2602 1 year, 9 months ago

Selected Answer: A

AWS Shield is a managed service that provides protection against Distributed Denial of Service (DDoS) attacks for applications running on AWS. AWS Shield Standard is automatically enabled to all AWS customers at no additional cost. AWS Shield Advanced is an optional paid service. AWS Shield Advanced provides additional protections against more sophisticated and larger attacks for your applications running on Amazon Elastic Compute Cloud (EC2), Elastic Load Balancing (ELB), Amazon CloudFront, AWS Global Accelerator, and Route 53.

upvoted 4 times

[Removed] 1 year, 9 months ago

Selected Answer: A

aaaaa

accelator can not be attached to shield

upvoted 2 times

[Removed] 1 year, 9 months ago

bbbbbbbbb

upvoted 1 times

enzomv 1 year, 8 months ago

Your origin servers can be Amazon Simple Storage Service (S3), Amazon EC2, Elastic Load Balancing, or a custom server outside of AWS. You can also enable AWS Shield Advanced directly on Elastic Load Balancing or Amazon EC2 in the following AWS Regions - Northern Virginia, Ohio, Oregon, Northern California, Montreal, São Paulo, Ireland, Frankfurt, London, Paris, Stockholm, Singapore, Tokyo, Sydney, Seoul, Mumbai, Milan, and Cape Town.

My answer is B

upvoted 2 times

pentium75 11 months, 2 weeks ago

You CAN enable Shield Advanced directly on EC2. You CAN also expose EC2 instances directly to the Internet. But in this case, what is exposed to the Internet (= where DDoS attacks could land) is the Global Accelerator, not your EC2 instances.

upvoted 3 times

enzomv 1 year, 8 months ago

<https://docs.aws.amazon.com/waf/latest/developerguide/ddos-event-mitigation-logic-gax.html>

Sorry I meant A

upvoted 3 times

ktulu2602 1 year, 9 months ago

Yes it can:

AWS Shield is a managed service that provides protection against Distributed Denial of Service (DDoS) attacks for applications running on AWS. AWS Shield Standard is automatically enabled to all AWS customers at no additional cost. AWS Shield Advanced is an optional paid service. AWS Shield Advanced provides additional protections against more sophisticated and larger attacks for your applications running on Amazon Elastic Compute Cloud (EC2), Elastic Load Balancing (ELB), Amazon CloudFront, AWS Global Accelerator, and Route 53.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #397

Topic 1

An ecommerce company needs to run a scheduled daily job to aggregate and filter sales records for analytics. The company stores the sales records in an Amazon S3 bucket. Each object can be up to 10 GB in size. Based on the number of sales events, the job can take up to an hour to complete. The CPU and memory usage of the job are constant and are known in advance.

A solutions architect needs to minimize the amount of operational effort that is needed for the job to run.

Which solution meets these requirements?

- A. Create an AWS Lambda function that has an Amazon EventBridge notification. Schedule the EventBridge event to run once a day.
- B. Create an AWS Lambda function. Create an Amazon API Gateway HTTP API, and integrate the API with the function. Create an Amazon EventBridge scheduled event that calls the API and invokes the function.
- C. Create an Amazon Elastic Container Service (Amazon ECS) cluster with an AWS Fargate launch type. Create an Amazon EventBridge scheduled event that launches an ECS task on the cluster to run the job. **Most Voted**
- D. Create an Amazon Elastic Container Service (Amazon ECS) cluster with an Amazon EC2 launch type and an Auto Scaling group with at least one EC2 instance. Create an Amazon EventBridge scheduled event that launches an ECS task on the cluster to run the job.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**ktulu2602** **Highly Voted** 1 year, 9 months ago**Selected Answer: C**

The requirement is to run a daily scheduled job to aggregate and filter sales records for analytics in the most efficient way possible. Based on the requirement, we can eliminate option A and B since they use AWS Lambda which has a limit of 15 minutes of execution time, which may not be sufficient for a job that can take up to an hour to complete.

Between options C and D, option C is the better choice since it uses AWS Fargate which is a serverless compute engine for containers that eliminates the need to manage the underlying EC2 instances, making it a low operational effort solution. Additionally, Fargate also provides instant scale-up and scale-down capabilities to run the scheduled job as per the requirement.

requirement.

Therefore, the correct answer is:

C. Create an Amazon Elastic Container Service (Amazon ECS) cluster with an AWS Fargate launch type. Create an Amazon EventBridge scheduled event that launches an ECS task on the cluster to run the job.

upvoted 27 times

imvb88 Highly Voted 1 year, 7 months ago

Selected Answer: C

"1-hour job" -> A, B out since max duration for Lambda is 15 min

Between C and D, "minimize operational effort" means Fargate -> C

upvoted 5 times

awsgeek75 Most Recent 11 months, 1 week ago

Selected Answer: C

A&B are out due to Lambda 15 min limits

C is less operationally complex than D so C is the right answer. Fargate is managed ECS cluster whereas EC2 based ECS will require more config overhead.

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

C. Create an Amazon Elastic Container Service (Amazon ECS) cluster with an AWS Fargate launch type. Create an Amazon EventBridge scheduled event that launches an ECS task on the cluster to run the job

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

The best option is C.

'The job can take up to an hour to complete' rules out lambda functions as they only execute up to 15 mins. Hence option A and B are out.

'The CPU and memory usage of the job are constant and are known in advance' rules out the need for autoscaling. Hence option D is out.

upvoted 4 times

klayytech 1 year, 8 months ago

Selected Answer: C

The solution that meets the requirements with the least operational overhead is to create a **Regional AWS WAF web ACL** with a rate-based rule and associate the web ACL with the API Gateway stage. This solution will protect the application from HTTP flood attacks by monitoring incoming requests and blocking requests from IP addresses that exceed the predefined rate.

Amazon CloudFront distribution with Lambda@Edge in front of the API Gateway Regional API endpoint is also a good solution but it requires more operational overhead than the previous solution.

Using Amazon CloudWatch metrics to monitor the Count metric and alerting the security team when the predefined rate is reached is not a solution that can protect against HTTP flood attacks.

Creating an Amazon CloudFront distribution in front of the API Gateway Regional API endpoint with a maximum TTL of 24 hours is not a solution that can protect against HTTP flood attacks.

upvoted 1 times

klayytech 1 year, 8 months ago

Selected Answer: C

The solution that meets these requirements is C. Create an Amazon Elastic Container Service (Amazon ECS) cluster with an AWS Fargate launch type. Create an Amazon EventBridge scheduled event that launches an ECS task on the cluster to run the job. This solution will minimize the amount of operational effort that is needed for the job to run.

AWS Lambda which has a limit of 15 minutes of execution time,

upvoted 2 times