Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #467 *Topic 1*

A company uses AWS Organizations. A member account has purchased a Compute Savings Plan. Because of changes in the workloads inside the member account, the account no longer receives the full benefit of the Compute Savings Plan commitment. The company uses less than 50% of its purchased compute power.

A. Turn on discount sharing from the Billing Preferences section of the account console in the member account that purchased the Compute Savings Plan.

B. Turn on discount sharing from the Billing Preferences section of the account console in the company's Organizations management account. **Most Voted**

C. Migrate additional compute workloads from another AWS account to the account that has the Compute Savings Plan.

D. Sell the excess Savings Plan commitment in the Reserved Instance Marketplace.

Correct Answer: *B*

Community vote distribution

B (74%) D (26%)

Comments

norris81 Highly Voted 2 years ago

Selected Answer: B

https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html

Sign in to the AWS Management Console and open the AWS Billing console at https://console.aws.amazon.com/billing/

.

Note

Ensure you're logged in to the management account of your AWS Organizations.

upvoted 11 times

baba365 Highly Voted 1 year, 8 months ago

So what exactly is the question?

upvoted 7 times

awsgeek75 1 year, 4 months ago

It's an English test on complete the sentence...

upvoted 4 times

pentium75 1 year, 5 months ago

What to do

upvoted 2 times

jingen11 Most Recent 7 months, 3 weeks ago

Selected Answer: B

I dont think there is a way for us to sell excess Savings plan. Only Selling Reserved instances is possible in marketplace

upvoted 2 times

jingen11 7 months, 3 weeks ago

https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html#ri-turn-on-process

upvoted 2 times

Stranko 1 year, 3 months ago

Selected Answer: D

I'd go with D, due to "The company uses less than 50% of its purchased compute power". Like, why are you sharing it between other accounts of the company, if the company itself doesn't need it? If you provisioned too much you can sell the overprovisioned capacity on the market. I'd understand B if it was about the account using about 50% of the plan and other accounts running similar workloads, but no such thing is stated.

upvoted 2 times

NayeraB 1 year, 3 months ago

Option E, Take it out of the salary of the guy who made the decision to purchase an entire compute plan without studying the company's needs.

upvoted 6 times

mr123dd 1 year, 5 months ago

Selected Answer: D

in the question, it does not clarify then number of accounts the company has, if they only has one account, I think it is D,

upvoted 1 times

Mujahid_1 1 year, 5 months ago

what are you guys doing
this section is for discussion not for copy paste

upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: B

B, it's a generic Compute Savings Plan that can be used for compute workloads in the other accounts.

A doesn't work, discount sharing must be enabled for all accounts (at least for those that provide and share the discounts).

C is not possible, there's a reason why the workloads are in different accounts.

D would be a last resort if there wouldn't be any other workloads in the own organization, but here are.

upvoted 4 times

michalf84 1 year, 8 months ago

Selected Answer: D

I saw similar question in older exam one can sell on the market unused capacity

upvoted 1 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

B. Turn on discount sharing from the Billing Preferences section of the account console in the company's Organizations

management account
upvoted 3 times

Lx016 1 year, 9 months ago

Bro, no need to copy paste the answer that is already written. Need an explanation, I see that you just copy pasting the potential answers without any explanation in each discussion.
upvoted 26 times

live_reply_developers 1 year, 11 months ago

Selected Answer: D

"For example, you might want to sell Reserved Instances after moving instances to a new AWS Region, changing to a new instance type, ending projects before the term expiration, when your business needs change, or if you have unneeded capacity."

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ri-market-general.html
upvoted 2 times

pentium75 1 year, 5 months ago

D would make sense if the company wouldn't have other accounts with workloads. Or if it would be EC2 Savings Plans that would not match the instance types in other accounts. But it's a generic Compute Savings Plan that surely can be used in another account. Thus B.
upvoted 2 times

awsgeek75 1 year, 4 months ago

I am also confused between B and D as the last part of the question "The company uses less than 50% of its purchased compute power." could imply that the whole company (not just this member account) only uses 50% of the computer power. If they said the member account only uses 50% then it would be clear cut B.
upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: B

answer is B.

https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html#:~:text=choose%20Save.-,Turning%20on%20shared%20reserved%20instances%20and%20Savings%20Plans%20discounts,-You%20can%20use
upvoted 2 times

Felix_br 2 years ago

Selected Answer: D

The company uses less than 50% of its purchased compute power.
For this reason i believe D is the best solution : https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ri-market-general.html
upvoted 3 times

Abrar2022 2 years ago

The company Organization's management account can turn on/off shared reserved instances.
upvoted 2 times

cloudenthusiast 2 years ago

Selected Answer: B

To summarize, option C (Migrate additional compute workloads from another AWS account to the account that has the Compute Savings Plan) is a valid solution to address the underutilization of the Compute Savings Plan. However, it involves workload migration and may require careful planning and coordination. Consider the feasibility and impact of migrating workloads before implementing this solution.
upvoted 3 times

EA100 2 years ago

Answer - C
If a member account within AWS Organizations has purchased a Compute Savings Plan
upvoted 1 times

EA100 2 years ago

Asnwer - C
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #468 *Topic 1*

A company is developing a microservices application that will provide a search catalog for customers. The company must use REST APIs to present the frontend of the application to users. The REST APIs must access the backend services that the company hosts in containers in private VPC subnets.

Which solution will meet these requirements?

A. Design a WebSocket API by using Amazon API Gateway. Host the application in Amazon Elastic Container Service (Amazon ECS) in a private subnet. Create a private VPC link for API Gateway to access Amazon ECS.

B. Design a REST API by using Amazon API Gateway. Host the application in Amazon Elastic Container Service (Amazon ECS) in a private subnet. Create a private VPC link for API Gateway to access Amazon ECS. **Most Voted**

C. Design a WebSocket API by using Amazon API Gateway. Host the application in Amazon Elastic Container Service (Amazon ECS) in a private subnet. Create a security group for API Gateway to access Amazon ECS.

D. Design a REST API by using Amazon API Gateway. Host the application in Amazon Elastic Container Service (Amazon ECS) in a private subnet. Create a security group for API Gateway to access Amazon ECS.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

cloudenthusiast Highly Voted 2 years ago

Selected Answer: B

REST API with Amazon API Gateway: REST APIs are the appropriate choice for providing the frontend of the microservices application. Amazon API Gateway allows you to design, deploy, and manage REST APIs at scale.

Amazon ECS in a Private Subnet: Hosting the application in Amazon ECS in a private subnet ensures that the containers are securely deployed within the VPC and not directly exposed to the public internet.

Private VPC Link: To enable the REST API in API Gateway to access the backend services hosted in Amazon ECS, you can create a private VPC link. This establishes a private network connection between the API Gateway and ECS containers, allowing secure communication without traversing the public internet.

upvoted 14 times

MNotABot Highly Voted 1 year, 11 months ago

Question itself says: "The company must use REST APIs", hence WebSocket APIs are not applicable and such options are eliminated straight away.
upvoted 9 times

MatAlves Most Recent 8 months, 4 weeks ago

Selected Answer: B

"VPC links enable you to create private integrations that connect your HTTP API routes to private resources in a VPC, such as Application Load Balancers or Amazon ECS container-based applications."
upvoted 3 times

freedafeng 10 months, 3 weeks ago

I think the connection should be from the application to the ECS in the private VPC, instead of from the API Gateway to the ECS in the private VPC. API Gateway only needs to connect to the application...
upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

AC are wrong as they are not REST API
D, you don't make SG for API Gateway to EC2, you have to make a VPC Link. More details at
https://docs.aws.amazon.com/apigateway/latest/developerguide/http-api-vpc-links.html
upvoted 5 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

To allow the REST APIs to securely access the backend, a private VPC link should be created from API Gateway to the ECS containers. A private VPC link provides private connectivity between API Gateway and the VPC without using public IP addresses or requiring an internet gateway/NAT
upvoted 5 times

Axeashes 1 year, 11 months ago

Selected Answer: B

https://docs.aws.amazon.com/apigateway/latest/developerguide/http-api-private-integration.html
upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: B

A VPC link is a resource in Amazon API Gateway that allows for connecting API routes to private resources inside a VPC.
upvoted 3 times

samehpalass 1 year, 11 months ago

B is the right choice
upvoted 2 times

Yadav_Sanjay 1 year, 11 months ago

Why Not D
upvoted 3 times

potomac 1 year, 7 months ago

A security group acts as a firewall for associated EC2 instances, controlling both inbound and outbound traffic at the instance level.
upvoted 2 times

nosense 2 years ago

Selected Answer: B

b is right, bcs vpc link provided security connection
upvoted 4 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #469 *Topic 1*

A company stores raw collected data in an Amazon S3 bucket. The data is used for several types of analytics on behalf of the company's customers. The type of analytics requested determines the access pattern on the S3 objects.

The company cannot predict or control the access pattern. The company wants to reduce its S3 costs.

Which solution will meet these requirements?

A. Use S3 replication to transition infrequently accessed objects to S3 Standard-Infrequent Access (S3 Standard-IA)

B. Use S3 Lifecycle rules to transition objects from S3 Standard to Standard-Infrequent Access (S3 Standard-IA)

C. Use S3 Lifecycle rules to transition objects from S3 Standard to S3 Intelligent-Tiering **Most Voted**

D. Use S3 Inventory to identify and transition objects that have not been accessed from S3 Standard to S3 Intelligent-Tiering

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

nosense Highly Voted 1 year ago

Selected Answer: C

S3 Inventory can't to move files to another class

upvoted 8 times

Murtadhaceit Highly Voted 6 months ago

Selected Answer: C

Unpredictable access pattern = Intelligent-Tiering.

upvoted 5 times

Guru4Cloud Most Recent 9 months, 2 weeks ago

Selected Answer: C

C. Use S3 Lifecycle rules to transition objects from S3 Standard to S3 Intelligent-Tiering

upvoted 4 times

TariqKipkemei 11 months, 3 weeks ago

Selected Answer: C

Cannot predict access pattern = S3 Intelligent-Tiering.

upvoted 4 times

Efren 1 year ago

Selected Answer: C

Not known patterns, Intelligent Tier

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #470 *Topic 1*

A company has applications hosted on Amazon EC2 instances with IPv6 addresses. The applications must initiate communications with other external applications using the internet. However the company's security policy states that any external service cannot initiate a connection to the EC2 instances.

What should a solutions architect recommend to resolve this issue?

A. Create a NAT gateway and make it the destination of the subnet's route table

B. Create an internet gateway and make it the destination of the subnet's route table

C. Create a virtual private gateway and make it the destination of the subnet's route table

D. Create an egress-only internet gateway and make it the destination of the subnet's route table **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

wRhlH Highly Voted 1 year, 11 months ago

For exam,
egress-only internet gateway: IPv6
NAT gateway: IPv4
upvoted 54 times

MatAlves 8 months, 4 weeks ago

Good stuff.

"An egress-only internet gateway is for use with IPv6 traffic only. To enable outbound-only internet communication over IPv4, use a NAT gateway instead."

https://docs.aws.amazon.com/vpc/latest/userguide/egress-only-internet-gateway.html
upvoted 2 times

b82faaf 1 year, 5 months ago

This is very helpful, thanks.

upvoted 4 times

RDM10 1 year, 8 months ago

thanks a lot

upvoted 4 times

cloudenthusiast Highly Voted 2 years ago

Selected Answer: D

An egress-only internet gateway (EIGW) is specifically designed for IPv6-only VPCs and provides outbound IPv6 internet access while blocking inbound IPv6 traffic. It satisfies the requirement of preventing external services from initiating connections to the EC2 instances while allowing the instances to initiate outbound communications.

upvoted 9 times

cloudenthusiast 2 years ago

Since the company's security policy explicitly states that external services cannot initiate connections to the EC2 instances, using a NAT gateway (option A) would not be suitable. A NAT gateway allows outbound connections from private subnets to the internet, but it does not restrict inbound connections from external sources.

upvoted 5 times

pentium75 1 year, 5 months ago

"A NAT gateway ... does not restrict inbound connections from external sources." Actually it does, but only for IPv4.

upvoted 2 times

[Removed] 2 years ago

Enable outbound IPv6 traffic using an egress-only internet gateway
https://docs.aws.amazon.com/vpc/latest/userguide/egress-only-internet-gateway.html

upvoted 3 times

MatAlves Most Recent 8 months, 4 weeks ago

Selected Answer: D

"An egress-only internet gateway is for use with IPv6 traffic only. To enable outbound-only internet communication over IPv4, use a NAT gateway instead."

https://docs.aws.amazon.com/vpc/latest/userguide/egress-only-internet-gateway.html

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

D. Create an egress-only internet gateway and make it the destination of the subnet's route table

upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: D

Outbound traffic only = Create an egress-only internet gateway and make it the destination of the subnet's route table

upvoted 2 times

radev 2 years ago

Selected Answer: D

Egress-Only internet Gateway

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #471 *Topic 1*

A company is creating an application that runs on containers in a VPC. The application stores and accesses data in an Amazon S3 bucket. During the development phase, the application will store and access 1 TB of data in Amazon S3 each day. The company wants to minimize costs and wants to prevent traffic from traversing the internet whenever possible.

Which solution will meet these requirements?

A. Enable S3 Intelligent-Tiering for the S3 bucket

B. Enable S3 Transfer Acceleration for the S3 bucket

C. Create a gateway VPC endpoint for Amazon S3. Associate this endpoint with all route tables in the VPC **Most Voted**

D. Create an interface endpoint for Amazon S3 in the VPC. Associate this endpoint with all route tables in the VPC

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

litos168 Highly Voted 1 year, 4 months ago

Amazon S3 supports both gateway endpoints and interface endpoints. With a gateway endpoint, you can access Amazon S3 from your VPC, without requiring an internet gateway or NAT device for your VPC, and with no additional cost. However, gateway endpoints do not allow access from on-premises networks, from peered VPCs in other AWS Regions, or through a transit gateway. For those scenarios, you must use an interface endpoint, which is available for an additional cost.

upvoted 14 times

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: C

Gateway VPC Endpoint: A gateway VPC endpoint enables private connectivity between a VPC and Amazon S3. It allows direct access to Amazon S3 without the need for internet gateways, NAT devices, VPN connections, or AWS Direct Connect.

Minimize Internet Traffic: By creating a gateway VPC endpoint for Amazon S3 and associating it with all route tables in the VPC, the traffic between the VPC and Amazon S3 will be kept within the AWS network. This helps in minimizing data transfer costs and prevents the need for traffic to traverse the internet.

Cost-Effective: With a gateway VPC endpoint, the data transfer between the application running in the VPC and the S3 bucket stays within the AWS network, reducing the need for data transfer across the internet. This can result in cost savings, especially

when dealing with large amounts of data.
upvoted 8 times

cloudenthusiast 1 year, 6 months ago

Option B (Enable S3 Transfer Acceleration for the S3 bucket) is a feature that uses the CloudFront global network to accelerate data transfers to and from Amazon S3. While it can improve data transfer speed, it still involves traffic traversing the internet and doesn't directly address the goal of minimizing costs and preventing internet traffic whenever possible.
upvoted 3 times

awsgeek75 Most Recent 10 months, 3 weeks ago

Selected Answer: C

https://aws.amazon.com/blogs/architecture/choosing-your-vpc-endpoint-strategy-for-amazon-s3/

A: Storage cost is not described as an issue here
B: Tx Accelerator is for external (global user) traffic acceleration
D: Interface endpoint is on-prem to S3
C: gateway VPC is specifically for S3 to AWS resources
upvoted 4 times

dkw2342 9 months ago

Interface endpoints are not exclusively for on-prem to S3.

The only reason why option D is wrong is because "Associate this endpoint with all route tables in the VPC" makes no sense.
upvoted 1 times

bsbs1234 1 year, 2 months ago

I think both C&D will works.
But D will have extra cost. So C is correct.
upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

C. Create a gateway VPC endpoint for Amazon S3. Associate this endpoint with all route tables in the VPC
upvoted 2 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: C

Prevent traffic from traversing the internet = Gateway VPC endpoint for S3.
upvoted 3 times

Anmol_1010 1 year, 6 months ago

Key word transversing to internet
upvoted 2 times

Efren 1 year, 6 months ago

Selected Answer: C

Gateway endpoint for S3
upvoted 3 times

nosense 1 year, 6 months ago

Selected Answer: C

vpc endpoint for s3
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #472 *Topic 1*

A company has a mobile chat application with a data store based in Amazon DynamoDB. Users would like new messages to be read with as little latency as possible. A solutions architect needs to design an optimal solution that requires minimal application changes.

Which method should the solutions architect select?

A. Configure Amazon DynamoDB Accelerator (DAX) for the new messages table. Update the code to use the DAX endpoint. **Most Voted**

B. Add DynamoDB read replicas to handle the increased read load. Update the application to point to the read endpoint for the read replicas.

C. Double the number of read capacity units for the new messages table in DynamoDB. Continue to use the existing DynamoDB endpoint.

D. Add an Amazon ElastiCache for Redis cache to the application stack. Update the application to point to the Redis cache endpoint instead of DynamoDB.

Correct Answer: *A*

Community vote distribution

A (95%) C (5%)

Comments

pentium75 11 months, 1 week ago

Selected Answer: A

B and C do not reduce latency. D would reduce latency but require significant application changes.

upvoted 2 times

wizcloudifa 7 months, 3 weeks ago

D would not reduce latency, all the messages are new, they wont be stored in cache and they are unique messages(dynamic content) so Elasticache would be pointless here, its more preferable for static content/frequently accessed content.... A makes perfect sense

upvoted 3 times

Cyberkayu 11 months, 3 weeks ago

Selected Answer: C

0 code change @C

ABD. In memory cache, read replica, elasticache. Chat application and content is dynamic, cache will still pull data from prod database

upvoted 1 times

pentium75 11 months, 1 week ago

C has 0 codes changes but doesn't address the issue.

upvoted 6 times

danielmakita 1 year, 1 month ago

Would go for A.
Minimal application changes != No application changes

upvoted 2 times

thanhnv142 1 year, 1 month ago

"requires minimal application changes" - Do not choose A because it requires updates of codes.

upvoted 1 times

thanhnv142 1 year, 1 month ago

C is correct
A, B and D all require code changes to the app.

upvoted 1 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

A. Configure Amazon DynamoDB Accelerator (DAX) for the new messages table. Update the code to use the DAX endpoint.

upvoted 2 times

haoAWS 1 year, 5 months ago

Selected Answer: A

Read replica does improve the read speed, but it cannot improve the latency because there is always latency between replicas. So A works and B not work.

upvoted 2 times

mattcl 1 year, 5 months ago

C , "requires minimal application changes"

upvoted 1 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: A

little latency = Amazon DynamoDB Accelerator (DAX) .

upvoted 4 times

DrWatson 1 year, 6 months ago

Selected Answer: A

I go with A https://aws.amazon.com/blogs/mobile/building-a-full-stack-chat-application-with-aws-and-nextjs/ but I have some doubts about this https://aws.amazon.com/blogs/database/how-to-build-a-chat-application-with-amazon-elasticache-for-redis/

upvoted 3 times

cloudenthusiast 1 year, 6 months ago

Selected Answer: A

Amazon DynamoDB Accelerator (DAX): DAX is an in-memory cache for DynamoDB that provides low-latency access to frequently accessed data. By configuring DAX for the new messages table, read requests for the table will be served from the DAX cache, significantly reducing the latency.

Minimal Application Changes: With DAX, the application code can be updated to use the DAX endpoint instead of the standard DynamoDB endpoint. This change is relatively minimal and does not require extensive modifications to the application's data access logic.

Low Latency: DAX caches frequently accessed data in memory, allowing subsequent read requests for the same data to be served with minimal latency. This ensures that new messages can be read by users with minimal delay.

upvoted 4 times

cloudenthusiast 1 year, 6 months ago

Option B (Add DynamoDB read replicas) involves creating read replicas to handle the increased read load, but it may not directly address the requirement of minimizing latency for new message reads.

upvoted 1 times

Efren 1 year, 6 months ago

Tricky one, in doubt also with B, read replicas.

upvoted 1 times

awsgeek75 11 months ago

Yes it's tricky but least code changes is the tie breaker. DAX has zero code changes.

upvoted 2 times

nosense 1 year, 6 months ago

Selected Answer: A

a is valid

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #473 *Topic 1*

A company hosts a website on Amazon EC2 instances behind an Application Load Balancer (ALB). The website serves static content. Website traffic is increasing, and the company is concerned about a potential increase in cost.

A. Create an Amazon CloudFront distribution to cache state files at edge locations **Most Voted**

B. Create an Amazon ElastiCache cluster. Connect the ALB to the ElastiCache cluster to serve cached files

C. Create an AWS WAF web ACL and associate it with the ALB. Add a rule to the web ACL to cache static files

D. Create a second ALB in an alternative AWS Region. Route user traffic to the closest Region to minimize data transfer costs

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: A

Amazon CloudFront: CloudFront is a content delivery network (CDN) service that caches content at edge locations worldwide. By creating a CloudFront distribution, static content from the website can be cached at edge locations, reducing the load on the EC2 instances and improving the overall performance.

Caching Static Files: Since the website serves static content, caching these files at CloudFront edge locations can significantly reduce the number of requests forwarded to the EC2 instances. This helps to lower the overall cost by offloading traffic from the instances and reducing the data transfer costs.

upvoted 6 times

awsgeek75 Highly Voted 11 months ago

Selected Answer: A

The problem with this question is that no sane AWS architect will chose any of these options and go for S3 caching. But given the choices, A is the only one which will solve the problem within reasonable cost.

upvoted 5 times

Guru4Cloud Most Recent 1 year, 3 months ago

Selected Answer: A

A. Create an Amazon CloudFront distribution to cache state files at edge locations
upvoted 3 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: A

Serves static content = Amazon CloudFront distribution.
upvoted 3 times

nosense 1 year, 6 months ago

Selected Answer: A

a for me
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #474 *Topic 1*

A company has multiple VPCs across AWS Regions to support and run workloads that are isolated from workloads in other Regions. Because of a recent application launch requirement, the company's VPCs must communicate with all other VPCs across all Regions.

Which solution will meet these requirements with the LEAST amount of administrative effort?

A. Use VPC peering to manage VPC communication in a single Region. Use VPC peering across Regions to manage VPC communications.

B. Use AWS Direct Connect gateways across all Regions to connect VPCs across regions and manage VPC communications.

C. Use AWS Transit Gateway to manage VPC communication in a single Region and Transit Gateway peering across Regions to manage VPC communications. Most Voted

D. Use AWS PrivateLink across all Regions to connect VPCs across Regions and manage VPC communications

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

Felix_br Highly Voted 1 year, 6 months ago

The correct answer is: C. Use AWS Transit Gateway to manage VPC communication in a single Region and Transit Gateway peering across Regions to manage VPC communications.

AWS Transit Gateway is a network hub that you can use to connect your VPCs and on-premises networks. It provides a single point of control for managing your network traffic, and it can help you to reduce the number of connections that you need to manage.

Transit Gateway peering allows you to connect two Transit Gateways in different Regions. This can help you to create a global network that spans multiple Regions.

To use Transit Gateway to manage VPC communication in a single Region, you would create a Transit Gateway in each Region. You would then attach your VPCs to the Transit Gateway.

To use Transit Gateway peering to manage VPC communication across Regions, you would create a Transit Gateway peering connection between the Transit Gateways in each Region.

upvoted 25 times

TariqKipkemei 1 year, 5 months ago

thank you for this comprehensive explanation

upvoted 3 times

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: C

AWS Transit Gateway: Transit Gateway is a highly scalable service that simplifies network connectivity between VPCs and on-premises networks. By using a Transit Gateway in a single Region, you can centralize VPC communication management and reduce administrative effort.

Transit Gateway Peering: Transit Gateway supports peering connections across AWS Regions, allowing you to establish connectivity between VPCs in different Regions without the need for complex VPC peering configurations. This simplifies the management of VPC communications across Regions.

upvoted 8 times

awsgeek75 Most Recent 11 months ago

Selected Answer: C

C is like a managed solution for A. A can work but with a lot of overhead (CIDR blocks uniqueness requirement). B and D are not the right products

upvoted 2 times

potomac 1 year, 1 month ago

Selected Answer: C

multiple regions + multiple VPCs --> Transit Gateway

upvoted 3 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: C

Definitely C.
Very well explained by @Felix_br

upvoted 2 times

omoakin 1 year, 6 months ago

Cccccccccccccccccccc
if you have services in multiple Regions, a Transit Gateway will allow you to access those services with a simpler network configuration.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #475 *Topic 1*

A company is designing a containerized application that will use Amazon Elastic Container Service (Amazon ECS). The application needs to access a shared file system that is highly durable and can recover data to another AWS Region with a recovery point objective (RPO) of 8 hours. The file system needs to provide a mount target m each Availability Zone within a Region.

A solutions architect wants to use AWS Backup to manage the replication to another Region.

Which solution will meet these requirements?

A. Amazon FSx for Windows File Server with a Multi-AZ deployment

B. Amazon FSx for NetApp ONTAP with a Multi-AZ deployment

C. Amazon Elastic File System (Amazon EFS) with the Standard storage class **Most Voted**

D. Amazon FSx for OpenZFS

Correct Answer: *C*

Community vote distribution

C (89%) B (11%)

Comments

elmogy Highly Voted 1 year, 6 months ago

Selected Answer: C

https://aws.amazon.com/efs/faq/
Q: What is Amazon EFS Replication?
EFS Replication can replicate your file system data to another Region or within the same Region without requiring additional infrastructure or a custom process. Amazon EFS Replication automatically and transparently replicates your data to a second file system in a Region or AZ of your choice. You can use the Amazon EFS console, AWS CLI, and APIs to activate replication on an existing file system. EFS Replication is continual and provides a recovery point objective (RPO) and a recovery time objective (RTO) of minutes, helping you meet your compliance and business continuity goals.

upvoted 13 times

wizcloudifa Most Recent 7 months, 3 weeks ago

Selected Answer: C

The key thing to notice here in the question is "with a recovery point objective (RPO) of 8 hours", as it is 8 hours of time it

The key thing to notice here in the question is "with a recovery point objective (RPO) of 8 hours", as it is 8 hours of time it recovery can be easily managed by EFS, no need to go for costlier and not built for this use-case(share file system) options like NetApp ONTAP(proprietary data cluster), OpenZFS(not a built in filesystem in AWS) or FSx for windows(file system for windows compatible workloads)

upvoted 3 times

awsgeek75 11 months ago

Selected Answer: C

A: ECS is not Windows File Server so won't work
B: ONTAP is proprietary data cluster completely unrelated to this question
D: OpenZFS needs a Linux kind of host for access. Not a built-in filesystem in AWS by default

upvoted 2 times

pentium75 11 months, 1 week ago

Selected Answer: C

"The file systemhttps://www.examtopics.com/exams/amazon/aws-certified-solutions-architect-associate-saa-c03/view/48/# needs to provide a mount target in each (!) Availability Zone within a Region", most regions have three AZs, but FSx Multi-AZ provides only nodes "spread across two AZs". While "or Amazon EFS file systems that use Regional storage classes [such as Standard], you can create a mount target in each Availability Zone in an AWS Region."

upvoted 3 times

pentium75 11 months, 1 week ago

Huh, comment has been scrambled a bit. Anyway

FSx Multi-AZ: Mount targets in two AZs
EFS Standard: Can create mount target in each AZ

upvoted 3 times

Goutham4981 1 year ago

Selected Answer: C

In the absence of this information, we can only make an assumption based on the provided requirements. The requirement for a shared file system that can recover data to another AWS Region with a recovery point objective (RPO) of 8 hours, and the need for a mount target in each Availability Zone within a Region, are all natively supported by Amazon EFS with the Standard storage class.

While Amazon FSx for NetApp ONTAP does provide shared file systems and supports both Windows and Linux, it does not natively support replication to another region through AWS Backup.

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

C. Amazon Elastic File System (Amazon EFS) with the Standard storage class

upvoted 2 times

cd93 1 year, 3 months ago

Selected Answer: B

B or C, but since question didn't mention operating system type, I guess we should go with B because it is more versatile (EFS supports Linux only), although ECS containers do support windows instances...

upvoted 2 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: C

Both option B and C will support this requirement.

https://aws.amazon.com/efs/faq/#:~:text=What%20is%20Amazon%20EFS%20Replication%3F

https://aws.amazon.com/fsx/netapp-ontap/faqs/#:~:text=How%20do%20I%20configure%20cross%2Dregion%20replication%20for%20the%20data%20in%20my%20file%20system%3F

upvoted 2 times

omoakin 1 year, 6 months ago

BBBBBBBBBBBBBBB
upvoted 1 times

[Removed] 1 year, 6 months ago

Both B and C are feasible.
Amazon FSx for NetApp ONTAP is just way overpriced for a backup storage solution. The keyword to look out for is sub milli seconds latency
In real life env, Amazon Elastic File System (Amazon EFS) with the Standard storage class is good enough.
upvoted 4 times

Anmol_1010 1 year, 6 months ago

Efs, can be mounted only in 1 region
So the answer is B
upvoted 3 times

Rob1L 1 year, 6 months ago

Selected Answer: C

C: EFS
upvoted 3 times

y0 1 year, 6 months ago

Selected Answer: C

AWS Backup can manage replication of EFS to another region as mentioned below
https://docs.aws.amazon.com/efs/latest/ug/awsbackup.html
upvoted 2 times

norris81 1 year, 6 months ago

https://aws.amazon.com/efs/faq/

During a disaster or fault within an AZ affecting all copies of your data, you might experience loss of data that has not been replicated using Amazon EFS Replication. EFS Replication is designed to meet a recovery point objective (RPO) and recovery time objective (RTO) of minutes. You can use AWS Backup to store additional copies of your file system data and restore them to a new file system in an AZ or Region of your choice. Amazon EFS file system backup data created and managed by AWS Backup is replicated to three AZs and is designed for 99.999999999% (11 nines) durability.
upvoted 2 times

nosense 1 year, 6 months ago

Amazon EFS is a scalable and durable elastic file system that can be used with Amazon ECS. However, it does not support replication to another AWS Region.
upvoted 1 times

fakrap 1 year, 6 months ago

To use EFS replication in a Region that is disabled by default, you must first opt in to the Region, so it does support.
upvoted 2 times

elmogy 1 year, 6 months ago

it does support replication to another AWS Region
check the same link you are replying to :/
https://aws.amazon.com/efs/faq/
Q: What is Amazon EFS Replication?
EFS Replication can replicate your file system data to another Region or within the same Region without requiring additional infrastructure or a custom process. Amazon EFS Replication automatically and transparently replicates your data to a second file system in a Region or AZ of your choice. You can use the Amazon EFS console, AWS CLI, and APIs to activate replication on an existing file system. EFS Replication is continual and provides a recovery point objective (RPO) and a recovery time objective (RTO) of minutes, helping you meet your compliance and business continuity goals.
upvoted 2 times

nosense 1 year, 6 months ago

Selected Answer: B

shared file system that is highly durable and can recover data
upvoted 2 times

Efren 1 year, 6 months ago

Why not EFS?
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #476 *Topic 1*

A company is expecting rapid growth in the near future. A solutions architect needs to configure existing users and grant permissions to new users on AWS. The solutions architect has decided to create IAM groups. The solutions architect will add the new users to IAM groups based on department.

Which additional action is the MOST secure way to grant permissions to the new users?

A. Apply service control policies (SCPs) to manage access permissions

B. Create IAM roles that have least privilege permission. Attach the roles to the IAM groups

C. Create an IAM policy that grants least privilege permission. Attach the policy to the IAM groups **Most Voted**

D. Create IAM roles. Associate the roles with a permissions boundary that defines the maximum permissions

Correct Answer: *C*

Community vote distribution

C (95%) B (5%)

Comments

Rob1L Highly Voted 2 years ago

Selected Answer: C

Option B is incorrect because IAM roles are not directly attached to IAM groups.

upvoted 9 times

RoroJ 2 years ago

IAM Roles can be attached to IAM Groups:
https://docs.aws.amazon.com/directoryservice/latest/admin-guide/assign_role.html

upvoted 4 times

antropaws 2 years ago

Read your own link: You can assign an existing IAM role to an AWS Directory Service user or group. Not to IAM groups.

upvoted 11 times

Efren Highly Voted 2 years ago

Selected Answer: C

Agreed with C

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_groups_manage_attach-policy.html

Attaching a policy to an IAM user group
upvoted 7 times

MatAlves Most Recent 8 months, 4 weeks ago

Selected Answer: C

"Manage access in AWS by creating policies and attaching them to IAM identities (users, groups of users, or roles) or AWS resources."

"An IAM role is an identity within your AWS account that has specific permissions. It's similar to an IAM user, but isn't associated with a specific person."

"IAM roles do not have any permanent credentials associated with them and are instead assumed by IAM users, AWS services, or applications that need temporary security credentials to access AWS resources"
upvoted 2 times

MatAlves 8 months, 4 weeks ago

https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies.html

https://docs.aws.amazon.com/IAM/latest/UserGuide/id.html

https://blog.awsfundamentals.com/aws-iam-roles-terms-concepts-and-examples
upvoted 2 times

zinabu 1 year, 2 months ago

create role=for resource like EC2 and lambda ....
create a Policy =for groups or user access policy for the resources like S3 bucket
upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: C

Not A or D because this is not about restricting maximum permissions, it is is about securely granting permissions

Not B because IAM roles are not attached to IAM groups.

C because IAM policies are attached to IAM groups.
upvoted 5 times

potomac 1 year, 7 months ago

Selected Answer: C

A is wrong
SCPs are mainly used along with AWS Organizations organizational units (OUs). SCPs do not replace IAM Policies such that they do not provide actual permissions. To perform an action, you would still need to grant appropriate IAM Policy permissions.
upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

Create an IAM policy that grants least privilege permission. Attach the policy to the IAM groups
upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: C

An IAM policy is an object in AWS that, when associated with an identity or resource, defines their permissions. Permissions in the policies determine whether a request is allowed or denied. You manage access in AWS by creating policies and attaching them to IAM identities (users, groups of users, or roles) or AWS resources.
So, option B will also work.
But Since I can only choose one, C would be it.
upvoted 3 times

MrAWSAssociate 1 year, 11 months ago

Selected Answer: C

You can attach up to 10 IAM policy for a 'user group'.
upvoted 2 times

antropaws 2 years ago

Selected Answer: C

C is the correct one.
upvoted 2 times

nosense 2 years ago

Selected Answer: B

should be b
upvoted 2 times

imazsyed 2 years ago

it should be C
upvoted 4 times

nosense 2 years ago

Option C is not as secure as option B because IAM policies are attached to individual users and cannot be used to manage permissions for groups of users.
upvoted 2 times

omoakin 2 years ago

IAM Roles manage who has access to your AWS resources, whereas IAM policies control their permissions. A Role with no Policy attached to it won't have to access any AWS resources. A Policy that is not attached to an IAM role is effectively unused.
upvoted 5 times

Clouddon 1 year, 9 months ago

https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies.html
upvoted 2 times

pentium75 1 year, 5 months ago

IAM roles are not attached to IAM groups.

IAM policies are attached to IAM roles, IAM groups or IAM users. IAM roles are used by services.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #477 *Topic 1*

A group requires permissions to list an Amazon S3 bucket and delete objects from that bucket. An administrator has created the following IAM policy to provide access to the bucket and applied that policy to the group. The group is not able to delete objects in the bucket. The company follows least-privilege access rules.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Action": [
                "s3:ListBucket",
                "s3:DeleteObject"
            ],
            "Resource": [
                "arn:aws:s3:::bucket-name"
            ],
            "Effect": "Allow"
        }
    ]
}
```

Which statement should a solutions architect add to the policy to correct bucket access?

A.
```
"Action": [
    "s3:*Object"
],
"Resource": [
    "arn:aws:s3:::bucket-name/*"
],
"Effect": "Allow"
```

B.
```
"Action": [
    "s3:*"
],
"Resource": [
    "arn:aws:s3:::bucket-name/*"
],
"Effect": "Allow"
```

```
"Action": [
```

```
        "s3:DeleteObject"
    ],
C.  "Resource": [
        "arn:aws:s3:::bucket-name*"
    ],
    "Effect": "Allow"
```

```
    "Action": [
        "s3:DeleteObject"
    ],
D.  "Resource": [
        "arn:aws:s3:::bucket-name/*"
    ],
    "Effect": "Allow"
```

Most Voted

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

Guru4Cloud Highly Voted 1 year, 3 months ago

Selected Answer: D

option B action is S3:*. this means all actions. The company follows least-privilege access rules. Hence option D

upvoted 6 times

nosense Highly Voted 1 year, 6 months ago

Selected Answer: D

d work

upvoted 5 times

Efren 1 year, 6 months ago

Agreed

upvoted 2 times

TariqKipkemei Most Recent 1 year, 5 months ago

Selected Answer: D

D is the answer

upvoted 4 times

AncaZalog 1 year, 5 months ago

what's the difference between B and D? on B the statements are just placed in another order

upvoted 1 times

TariqKipkemei 1 year, 5 months ago

option B action is S3:*. this means all actions. The company follows least-privilege access rules. Hence option D

upvoted 2 times

serepetru 1 year, 6 months ago

What is the difference between C and D?

upvoted 2 times

Ta_Les 1 year, 5 months ago

the "/" at the end of the last line on D

upvoted 6 times

sheilawu 6 months, 2 weeks ago

so annoying when you oversee this"/"
upvoted 1 times

Rob1L 1 year, 6 months ago

Selected Answer: D

D for sure
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #478 *Topic 1*

A law firm needs to share information with the public. The information includes hundreds of files that must be publicly readable. Modifications or deletions of the files by anyone before a designated future date are prohibited.

Which solution will meet these requirements in the MOST secure way?

A. Upload all files to an Amazon S3 bucket that is configured for static website hosting. Grant read-only IAM permissions to any AWS principals that access the S3 bucket until the designated date.

B. Create a new Amazon S3 bucket with S3 Versioning enabled. Use S3 Object Lock with a retention period in accordance with the designated date. Configure the S3 bucket for static website hosting. Set an S3 bucket policy to allow read-only access to the objects. Most Voted

C. Create a new Amazon S3 bucket with S3 Versioning enabled. Configure an event trigger to run an AWS Lambda function in case of object modification or deletion. Configure the Lambda function to replace the objects with the original versions from a private S3 bucket.

D. Upload all files to an Amazon S3 bucket that is configured for static website hosting. Select the folder that contains the files. Use S3 Object Lock with a retention period in accordance with the designated date. Grant read-only IAM permissions to any AWS principals that access the S3 bucket.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

Lin878 Highly Voted 11 months, 3 weeks ago

Selected Answer: B

Object Lock works only in buckets that have S3 Versioning enabled.
https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.html

upvoted 6 times

nosense Highly Voted 2 years ago

Selected Answer: B

Option A allows the files to be modified or deleted by anyone with read-only IAM permissions. Option C allows the files to be

modified or deleted by anyone who can trigger the AWS Lambda function.
Option D allows the files to be modified or deleted by anyone with read-only IAM permissions to the S3 bucket
upvoted 5 times

wizcloudifa 1 year, 1 month ago

no it doesnt, did you not notice this part: "S3 Object Lock with a retention period in accordance with the designated date", this part avoids deletion/modification of files
upvoted 2 times

eduvaldi Most Recent 3 days, 5 hours ago

Selected Answer: B

Versioning and Object Lock is necessary
upvoted 1 times

MatAlves 8 months, 4 weeks ago

Selected Answer: B

Versioning Enabled + Object Lock = B
upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: B

S3 bucket policy
upvoted 4 times

thanhnv142 1 year, 7 months ago

B is correct.
A doesnot have S3 object lock, but deletion is prohibited, which implies object lock
C does not have S3 as static web, but have to share the s3 with the public
D mentions files - but S3 manages objects, not file
upvoted 2 times

hydro143 1 year, 8 months ago

D?
Its like B, but also with read-only access limitations for anyone with IAM permissions. Also versioning in B doesn't help with anything.
upvoted 2 times

ManikRoy 1 year ago

Enabling versioning is a pre-requisite for object lock.
upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

Create a new Amazon S3 bucket with S3 Versioning enabled. Use S3 Object Lock with a retention period in accordance with the designated date. Configure the S3 bucket for static website hosting. Set an S3 bucket policy to allow read-only access to the objects.
upvoted 3 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: B

Create a new Amazon S3 bucket with S3 Versioning enabled. Use S3 Object Lock with a retention period in accordance with the designated date. Configure the S3 bucket for static website hosting. Set an S3 bucket policy to allow read-only access to the objects.
upvoted 4 times

antropaws 2 years ago

Selected Answer: B

Clearly B.

upvoted 3 times

dydzah 2 years ago

Selected Answer: B

https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.html

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #479 *Topic 1*

A company is making a prototype of the infrastructure for its new website by manually provisioning the necessary infrastructure. This infrastructure includes an Auto Scaling group, an Application Load Balancer and an Amazon RDS database. After the configuration has been thoroughly validated, the company wants the capability to immediately deploy the infrastructure for development and production use in two Availability Zones in an automated fashion.

What should a solutions architect recommend to meet these requirements?

A. Use AWS Systems Manager to replicate and provision the prototype infrastructure in two Availability Zones

B. Define the infrastructure as a template by using the prototype infrastructure as a guide. Deploy the infrastructure with AWS CloudFormation. Most Voted

C. Use AWS Config to record the inventory of resources that are used in the prototype infrastructure. Use AWS Config to deploy the prototype infrastructure into two Availability Zones.

D. Use AWS Elastic Beanstalk and configure it to use an automated reference to the prototype infrastructure to automatically deploy new environments in two Availability Zones.

Correct Answer: *B*

Community vote distribution

B (84%) A (16%)

Comments

Guru4Cloud Highly Voted 1 year, 9 months ago

Just Think Infrastructure as Code=== Cloud Formation

upvoted 9 times

nosense Highly Voted 2 years ago

Selected Answer: B

b obvious

upvoted 5 times

MatAlves Most Recent 8 months, 4 weeks ago

Selected Answer: A

The difference between CloudFormation and Beanstalk might be trick, but just for the exam think:

Cloudformation -> Infra as Code
Beanstalk -> deploy and manage applications
upvoted 5 times

LeonSauveterre 6 months ago

That's right. And I think you've voted for A by mistake. I think you wanna vote for B ;)
upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

A: Wrong product
C: Wrong product
D: EBS can only handle EC2 so RDS won't be replicated automatically
B: CloudFormation = IaaC
upvoted 4 times

capino 1 year, 9 months ago

Selected Answer: B

Just Think Infrastructure as Code=== Cloud Formation
upvoted 5 times

haoAWS 1 year, 11 months ago

Why D is not correct?
upvoted 2 times

Kiki_Pass 1 year, 10 months ago

I guess it's because Beanstalk is PaaS (platform as a service) while CloudFormation is IaC (infrastructure as code). The question emphasis more on infrastructure
upvoted 3 times

wRhlH 1 year, 11 months ago

I guess "TEMPLATE" leads to CloudFormation
upvoted 3 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: B

Infrastructure as code = AWS CloudFormation
upvoted 5 times

antropaws 2 years ago

Selected Answer: B

Clearly B.
upvoted 3 times

Felix_br 2 years ago

Selected Answer: B

AWS CloudFormation is a service that allows you to define and provision infrastructure as code. This means that you can create a template that describes the resources you want to create, and then use CloudFormation to deploy those resources in an automated fashion.

In this case, the solutions architect should define the infrastructure as a template by using the prototype infrastructure as a guide. The template should include resources for an Auto Scaling group, an Application Load Balancer, and an Amazon RDS database. Once the template is created, the solutions architect can use CloudFormation to deploy the infrastructure in two Availability Zones.
upvoted 4 times

omoakin 2 years ago

B
Define the infrastructure as a template by using the prototype infrastructure as a guide. Deploy the infrastructure with AWS CloudFormation

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #480 *Topic 1*

A business application is hosted on Amazon EC2 and uses Amazon S3 for encrypted object storage. The chief information security officer has directed that no application traffic between the two services should traverse the public internet.

Which capability should the solutions architect use to meet the compliance requirements?

A. AWS Key Management Service (AWS KMS)
B. VPC endpoint Most Voted
C. Private subnet
D. Virtual private gateway

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

cloudenthusiast Highly Voted 1 year ago

Selected Answer: B

A VPC endpoint enables you to privately access AWS services without requiring internet gateways, NAT gateways, VPN connections, or AWS Direct Connect connections. It allows you to connect your VPC directly to supported AWS services, such as Amazon S3, over a private connection within the AWS network.

By creating a VPC endpoint for Amazon S3, the traffic between your EC2 instances and S3 will stay within the AWS network and won't traverse the public internet. This provides a more secure and compliant solution, as the data transfer remains within the private network boundaries.

upvoted 10 times

Tendu Most Recent 3 months, 2 weeks ago

Selected Answer: B

Answer B -VPC Endpoint

upvoted 1 times

TariqKipkemei 11 months, 3 weeks ago

Selected Answer: B

Prevent traffic from traversing the internet = VPC endpoint for S3.
upvoted 4 times

antropaws 1 year ago

Selected Answer: B

B until proven contrary.
upvoted 3 times

handsonlabsaws 1 year ago

Selected Answer: B

B for sure
upvoted 3 times

Blingy 1 year ago

BBBBBBBBB
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #481 *Topic 1*

A company hosts a three-tier web application in the AWS Cloud. A Multi-AZAmazon RDS for MySQL server forms the database layer Amazon ElastiCache forms the cache layer. The company wants a caching strategy that adds or updates data in the cache when a customer adds an item to the database. The data in the cache must always match the data in the database.

Which solution will meet these requirements?

A. Implement the lazy loading caching strategy

B. Implement the write-through caching strategy **Most Voted**

C. Implement the adding TTL caching strategy

D. Implement the AWS AppConfig caching strategy

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: B

In the write-through caching strategy, when a customer adds or updates an item in the database, the application first writes the data to the database and then updates the cache with the same data. This ensures that the cache is always synchronized with the database, as every write operation triggers an update to the cache.

upvoted 30 times

cloudenthusiast 1 year, 6 months ago

Lazy loading caching strategy (option A) typically involves populating the cache only when data is requested, and it does not guarantee that the data in the cache always matches the data in the database.

Adding TTL (Time-to-Live) caching strategy (option C) involves setting an expiration time for cached data. It is useful for scenarios where the data can be considered valid for a specific period, but it does not guarantee that the data in the cache is always in sync with the database.

AWS AppConfig caching strategy (option D) is a service that helps you deploy and manage application configurations. It is not specifically designed for caching data synchronization between a database and cache layer.

upvoted 37 times

Kp88 1 year, 4 months ago

Great explanation , thanks

upvoted 3 times

zinabu Most Recent 8 months ago

write-through cashing strategy

upvoted 2 times

dikshya1233 10 months, 2 weeks ago

In exam

upvoted 4 times

awsgeek75 10 months, 3 weeks ago

Selected Answer: B

More helpful reading for why B is the answer:

https://docs.aws.amazon.com/AmazonElastiCache/latest/mem-ug/Strategies.html#Strategies.WriteThrough

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

B. Implement the write-through caching strategy

upvoted 4 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: B

The answer is definitely B.
I couldn't provide any more details than what has been shared by @cloudenthusiast.

upvoted 2 times

nosense 1 year, 6 months ago

Selected Answer: B

write-through caching strategy updates the cache at the same time as the database

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #482 *Topic 1*

A company wants to migrate 100 GB of historical data from an on-premises location to an Amazon S3 bucket. The company has a 100 megabits per second (Mbps) internet connection on premises. The company needs to encrypt the data in transit to the S3 bucket. The company will store new data directly in Amazon S3.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use the s3 sync command in the AWS CLI to move the data directly to an S3 bucket

B. Use AWS DataSync to migrate the data from the on-premises location to an S3 bucket **Most Voted**

C. Use AWS Snowball to move the data to an S3 bucket

D. Set up an IPsec VPN from the on-premises location to AWS. Use the s3 cp command in the AWS CLI to move the data directly to an S3 bucket

Correct Answer: *B*

Community vote distribution

B (62%) A (38%)

Comments

cloudenthusiast Highly Voted 2 years ago

Selected Answer: B

AWS DataSync is a fully managed data transfer service that simplifies and automates the process of moving data between on-premises storage and Amazon S3. It provides secure and efficient data transfer with built-in encryption, ensuring that the data is encrypted in transit.

By using AWS DataSync, the company can easily migrate the 100 GB of historical data from their on-premises location to an S3 bucket. DataSync will handle the encryption of data in transit and ensure secure transfer.

upvoted 13 times

luiscc Highly Voted 2 years ago

Selected Answer: B

Using DataSync, the company can easily migrate the 100 GB of historical data to an S3 bucket. DataSync will handle the encryption of data in transit, so the company does not need to set up a VPN or worry about managing encryption keys.

Option A, using the s3 sync command in the AWS CLI to move the data directly to an S3 bucket, would require more operational overhead as the company would need to manage the encryption of data in transit themselves. Option D, setting up operational overhead as the company would need to manage the encryption of data in transit themselves. Option D, setting up an IPsec VPN from the on-premises location to AWS, would also require more operational overhead and would be overkill for this scenario. Option C, using AWS Snowball, could work but would require more time and resources to order and set up the physical device.

upvoted 7 times

Dz1990 Most Recent 2 weeks, 1 day ago

Selected Answer: B

s3 sync Disadvantages:

Manual process = You handle errors, retries, progress monitoring
Basic tool = No optimization for large transfers
Network issues = Manual restart if connection drops
No built-in throttling = Could saturate bandwidth
Higher operational overhead = More hands-on management needed

upvoted 1 times

Manimgh 1 month ago

Selected Answer: B

encrypted = DataSync

upvoted 1 times

bora4motion 1 month, 1 week ago

Selected Answer: B

I will not use cli (A) to send 100GB to s3 from web. This means I will have to stay by computer? or use a different ec2 to login into it and initiate cli from there? hell no!!
B Datasync is the elegant way.

upvoted 1 times

dollykidurian 2 months, 3 weeks ago

My counter to all those, picking choice A, you will have to manually re-run the command if on-premise internet goes down however with datasync the service auto resume. (less operational overhead and less monitoring)

upvoted 1 times

zdi561 4 months ago

Selected Answer: B

You can also copy or synchronize data from on-premises to AWS manually with the AWS CLI, but doing so requires a fair bit of system tuning see https://aws.amazon.com/blogs/storage/migrating-and-managing-large-datasets-on-amazon-s3/

upvoted 1 times

FlyingHawk 6 months ago

Selected Answer: A

For A, it is more straightforward approach, with 100Mbps = 12.5MBs, 100 GB = 100 * 1024 MB = 102,400 MB, 102400 / 12.5 = 8192s = 2.3 hours, so if the network is stable, it should be less operational overhead than B which requires the installation of agent and network configuration, so compared to S3 sync, DataSync requires more upfront configuration but provides more robust, automated transfer capabilities.In an AWS certification exam, they typically expect you to choose the solution that provides the most straightforward, efficient approach with the least operational overhead while meeting all specified requirements.
For this scenario, they would likely expect you to choose A.

upvoted 1 times

FlyingHawk 6 months ago

In the context of AWS exam, AWS may expect we choose B as AWS DataSync is a managed service designed specifically for data migration tasks, A is more ad-hoc approach suit for small data-set. 100 GB will take 2.22 hours, although it is acceptable.

upvoted 2 times

LeonSauveterre 6 months ago

Selected Answer: A

Option A meets all the requirements with the LEAST operational overhead (while option B will introduce unnecessary complexity for a one-time transfer). The data will be encrypted in transit (via HTTPS), and a 100 Mbps connection is sufficient to transfer data of 100 GB.

If the task involves periodic transfers or error recovery, DataSync (option B) could be considered as an alternative.
upvoted 2 times

1e22522 10 months, 1 week ago

Selected Answer: B

Why would u u se the CLI
upvoted 2 times

bujuman 1 year, 2 months ago

Selected Answer: A

Assertions:
- needs to encrypt the data in transit to the S3 bucket.
- The company will store new data directly in Amazon S3.
Requirements:
- with the LEAST operational overhead
Even Though options A and B could do the job, option A requires VM maintenance because it is not a once-off migration (The company will store new data directly in Amazon S3)
NB: According to me, we must stuck to the question and avoid to interpret
upvoted 3 times

bujuman 1 year, 2 months ago

Erratum:
Assertions:
- needs to encrypt the data in transit to the S3 bucket.
- The company will store new data directly in Amazon S3.
Requirements:
- with the LEAST operational overhead
Even Though options A and B could do the job, option B requires VM maintenance because it is not a once-off migration (The company will store new data directly in Amazon S3)
NB: According to me, we must stuck to the question and avoid to interpret
upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: A

A - one single command, uses encryption automatically
B - Must install, configure and eventually decommission DataSync
C - Overkill
D - No need for VPN
upvoted 5 times

awsgeek75 1 year, 4 months ago

I agree, A is a million times simpler than B in terms of operational setup. AWS CLI is just one install on a server on client side and one command (literally) to sync the data.
upvoted 3 times

1rob 1 year, 5 months ago

Selected Answer: A

By default, all data transmitted from the client computer running the AWS CLI and AWS service endpoints is encrypted by sending everything through a HTTPS/TLS connection. You don't need to do anything to enable the use of HTTPS/TLS. It is always enabled unless you explicitly disable it for an individual command by using the --no-verify-ssl command line option. This is simpler compared to datasync, which will cost operational overhead to configure.
upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: B

storage data (including metadata) is encrypted in transit, but how it's encrypted throughout the transfer depends on your source and destination locations.
upvoted 2 times

thanhnv142 1 year, 7 months ago

B is correct to migrate
A is incorrect because is it only used to upload minor files (about a few GB) to AWS. 100 GB is not appropriate.
upvoted 2 times

awsgeek75 1 year, 4 months ago

There is no limitation on AWS CLI s3 sync command transfer size. Not that I can find in the docs.
https://awscli.amazonaws.com/v2/documentation/api/latest/reference/s3/sync.html

Happy to be corrected!
upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

Use AWS DataSync to migrate the data from the on-premises location to an S3 bucket
upvoted 4 times

HectorLeon2099 1 year, 10 months ago

Selected Answer: A

B is a good option but as the volume is not large and the speed is not bad, A requires less operational overhead
upvoted 5 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #483 *Topic 1*

A company containerized a Windows job that runs on .NET 6 Framework under a Windows container. The company wants to run this job in the AWS Cloud. The job runs every 10 minutes. The job's runtime varies between 1 minute and 3 minutes.

Which solution will meet these requirements MOST cost-effectively?

A. Create an AWS Lambda function based on the container image of the job. Configure Amazon EventBridge to invoke the function every 10 minutes.

B. Use AWS Batch to create a job that uses AWS Fargate resources. Configure the job scheduling to run every 10 minutes.

C. Use Amazon Elastic Container Service (Amazon ECS) on AWS Fargate to run the job. Create a scheduled task based on the container image of the job to run every 10 minutes. **Most Voted**

D. Use Amazon Elastic Container Service (Amazon ECS) on AWS Fargate to run the job. Create a standalone task based on the container image of the job. Use Windows task scheduler to run the job every
10 minutes.

Correct Answer: *C*

Community vote distribution

C (51%) B (29%) A (20%)

Comments

baba365 Highly Voted 1 year, 8 months ago

Lambda supports only Linux-based container images.

https://docs.aws.amazon.com/lambda/latest/dg/images-create.html

upvoted 16 times

awsgeek75 1 year, 5 months ago

Not really. Lambda supports .Net 6 directly: https://aws.amazon.com/blogs/compute/introducing-the-net-6-runtime-for-aws-lambda/

upvoted 7 times

AmrFawzy93 Highly Voted 2 years ago

Selected Answer: C

By using Amazon ECS on AWS Fargate, you can run the job in a containerized environment while benefiting from the serverless nature of Fargate, where you only pay for the resources used during the job's execution. Creating a scheduled task based on the container image of the job ensures that it runs every 10 minutes, meeting the required schedule. This solution provides flexibility, scalability, and cost-effectiveness.
upvoted 8 times

bora4motion 1 month, 1 week ago

why not lambda?
upvoted 1 times

zdi561 Most Recent 4 months ago

Selected Answer: A

Lambda is the best for cost , the job is kicked off every 10min and it only takes several min.Set 10min timeout , there is no cost for idling . Batch and ECS cost more for sure.
upvoted 1 times

zdi561 4 months, 2 weeks ago

Selected Answer: A

Lambda can be in .net container, and it is the cheaper than any other options.
upvoted 1 times

LeonSauveterre 6 months ago

Selected Answer: C

For B: Even though AWS Batch is free, I think it's more suitable for large-scale or relatively complex workloads, not lightweight periodic jobs. So B is not cost-effective (The cost depends on the compute resources you provision to execute the jobs, and maybe "cost-effective" somewhat involves less operational overhead because you have spent your money)

If you require queueing for task prioritization or processing a backlog of tasks, or work on data processing over datasets where Batch excels, AWS Batch would be better!
upvoted 1 times

muhammadahmer36 10 months, 1 week ago

A A A A A A
upvoted 2 times

zinabu 1 year, 2 months ago

selected answer : A
AWS Lambda now supports .NET 6 as both a managed runtime and a container base image
upvoted 5 times

xBUGx 1 year, 2 months ago

Selected Answer: A

https://aws.amazon.com/about-aws/whats-new/2022/02/aws-lambda-adds-support-net6/
upvoted 3 times

awsgeek75 1 year, 5 months ago

The question is weirdly phrased for .Net based containers. "A company containerized a Windows job that runs on .NET 6 Framework under a Windows container." This could mean that the job requires .Net 6 Framework OR it could mean the job requires Windows and .Net Framework 6. If the job is just based on .Net 6 then Lambda can run it. I am just a bit cautious about language because other parameters fall under Lambda. Question may have been wrongly quoted here.
upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: C

I guess this is an old question from before August 2023, when AWS Batch did not support Windows containers, while ECS already did since September 2021. Thus it would be C, though now B does also work. Since both Batch and ECS are free, we'd pay only for the Fargate resources (which are identical in both cases), now B and C would be correct.

A doesn't work because Lambda still does not support Windows containeres.
D doesn't make sense because the container would have to run 24/7

D doesn't make sense because the container would have to run 24/7
upvoted 7 times

ftaws 1 year, 5 months ago

I think that Batch with Fargate is more cheaper than ECS.
upvoted 1 times

pentium75 1 year, 5 months ago

Both Batch and ECS are free.
https://aws.amazon.com/de/ecs/pricing/
https://aws.amazon.com/de/batch/pricing/
upvoted 2 times

kt7 1 year, 6 months ago

Selected Answer: B

Batch supports fargate now
upvoted 6 times

ccmc 1 year, 7 months ago

Selected Answer: B

aws batch supports fargate
upvoted 3 times

deechean 1 year, 9 months ago

Selected Answer: C

C works. For A, the lambda support container image, but the container image much implement the Lambda Runtime API.
upvoted 1 times

markoniz 1 year, 8 months ago

Absolutely agree with this one ... Lambda do not support Windows container, on the other hand ECS is adequate solution
upvoted 3 times

Hades2231 1 year, 9 months ago

Selected Answer: B

As they support Batch on Fargate now (Aug 2023), the correct answer should be B?
upvoted 3 times

RDM10 1 year, 8 months ago

that's exactly my question too.
In one of the discussions, they same lambda is for jobs for 15 min. But for other question, they same batch is the best. I do not understand why we cant use batch?
upvoted 1 times

Smart 1 year, 9 months ago

Selected Answer: A

https://docs.aws.amazon.com/lambda/latest/dg/csharp-image.html#csharp-image-clients
upvoted 1 times

pentium75 1 year, 5 months ago

But it's clearly "a Windows job". Lambda does not support Windows containers. (.NET 6 could also run under Linux, but they'd need to modify the container in any case.)
upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

C is the most cost-effective solution for running a short-lived Windows container job on a schedule.

Using Amazon ECS scheduled tasks on Fargate eliminates the need to provision EC2 resources. You pay only for the duration the task runs.

Scheduled tasks handle scheduling the jobs and scaling resources automatically. This is lower cost than managing your own scaling via Lambda or Batch.

ECS also supports Windows containers natively unlike Lambda (option A).

Option D still requires provisioning and paying for full time EC2 resources to run a task scheduler even when tasks are not running.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #484 *Topic 1*

A company wants to move from many standalone AWS accounts to a consolidated, multi-account architecture. The company plans to create many new AWS accounts for different business units. The company needs to authenticate access to these AWS accounts by using a centralized corporate directory service.

Which combination of actions should a solutions architect recommend to meet these requirements? (Choose two.)

A. Create a new organization in AWS Organizations with all features turned on. Create the new AWS accounts in the organization. **Most Voted**

B. Set up an Amazon Cognito identity pool. Configure AWS IAM Identity Center (AWS Single Sign-On) to accept Amazon Cognito authentication.

C. Configure a service control policy (SCP) to manage the AWS accounts. Add AWS IAM Identity Center (AWS Single Sign-On) to AWS Directory Service.

D. Create a new organization in AWS Organizations. Configure the organization's authentication mechanism to use AWS Directory Service directly.

E. Set up AWS IAM Identity Center (AWS Single Sign-On) in the organization. Configure IAM Identity Center, and integrate it with the company's corporate directory service. **Most Voted**

Correct Answer: *AE*

Community vote distribution

AE (100%)

Comments

cloudenthusiast Highly Voted 1 year ago

Selected Answer: AE

A. By creating a new organization in AWS Organizations, you can establish a consolidated multi-account architecture. This allows you to create and manage multiple AWS accounts for different business units under a single organization.

E. Setting up AWS IAM Identity Center (AWS Single Sign-On) within the organization enables you to integrate it with the company's corporate directory service. This integration allows for centralized authentication, where users can sign in using their corporate credentials and access the AWS accounts within the organization.

Together, these actions create a centralized, multi-account architecture that leverages AWS Organizations for account management and AWS IAM Identity Center (AWS Single Sign-On) for authentication and access control.

management and AWS IAM Identity Center (AWS Single Sign-On) for authentication and access control.
upvoted 11 times

Guru4Cloud Most Recent 9 months, 2 weeks ago

Selected Answer: AE

A) Using AWS Organizations allows centralized management of multiple AWS accounts in a single organization. New accounts can easily be created within the organization.

E) Integrating AWS IAM Identity Center (AWS SSO) with the company's corporate directory enables federated single sign-on. Users can log in once to access accounts and resources across AWS.

Together, Organizations and IAM Identity Center provide consolidated management and authentication for multiple accounts using existing corporate credentials.
upvoted 3 times

samehpalass 11 months, 2 weeks ago

Selected Answer: AE

A:AWS Organization
E:Authentication because option C (SCP) for Authorization
upvoted 4 times

baba365 11 months ago

Ans: CD

'centralized corporate directory service' with new accounts in AWS Organizations
upvoted 2 times

TariqKipkemei 11 months, 3 weeks ago

Selected Answer: AE

Create a new organization in AWS Organizations with all features turned on. Create the new AWS accounts in the organization. Set up AWS IAM Identity Center (AWS Single Sign-On) in the organization. Configure IAM Identity Center, and integrate it with the company's corporate directory service.

AWS IAM Identity Center (successor to AWS Single Sign-On) helps you securely create or connect your workforce identities and manage their access centrally across AWS accounts and applications.

https://aws.amazon.com/iam/identity-center/#:~:text=AWS%20IAM%20Identity%20Center%20(successor%20to%20AWS%20Single%20Sign%2DOn)%20helps%20you%20securely%20create%20or%20connect%20your%20workforce%20identities%20and%20manage%20their%20access%20centrally%20across%20AWS%20accounts%20and%20applications.
upvoted 2 times

nosense 1 year ago

ae is right
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #485 *Topic 1*

A company is looking for a solution that can store video archives in AWS from old news footage. The company needs to minimize costs and will rarely need to restore these files. When the files are needed, they must be available in a maximum of five minutes.

What is the MOST cost-effective solution?

A. Store the video archives in Amazon S3 Glacier and use Expedited retrievals. **Most Voted**

B. Store the video archives in Amazon S3 Glacier and use Standard retrievals.

C. Store the video archives in Amazon S3 Standard-Infrequent Access (S3 Standard-IA).

D. Store the video archives in Amazon S3 One Zone-Infrequent Access (S3 One Zone-IA).

Correct Answer: *A*

Community vote distribution

A (97%) C

Comments

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: A

By choosing Expedited retrievals in Amazon S3 Glacier, you can reduce the retrieval time to minutes, making it suitable for scenarios where quick access is required. Expedited retrievals come with a higher cost per retrieval compared to standard retrievals but provide faster access to your archived data.

upvoted 12 times

nosense Highly Voted 1 year, 6 months ago

glacier expedited retrieval times of typically 1-5 minutes.

upvoted 5 times

wsdasdasdqwdaw 1 year, 1 month ago

Fully agree. Check here for evidences: https://aws.amazon.com/s3/storage-classes/glacier/#:~:text=S3%20Glacier%20Flexible%20Retrieval%20provides,amounts%20of%20data%20typically%20in

upvoted 2 times

Ravan Most Recent 9 months, 2 weeks ago

Selected Answer: A

The most cost-effective solution that also meets the requirement of having the files available within a maximum of five minutes when needed is:

A. Store the video archives in Amazon S3 Glacier and use Expedited retrievals.

Amazon S3 Glacier is designed for long-term storage of data archives, providing a highly durable and secure solution at a low cost. With Expedited retrievals, data can be retrieved within a few minutes, which meets the requirement of having the files available within five minutes when needed. This option provides the balance between cost-effectiveness and retrieval speed, making it the best choice for the company's needs.

upvoted 2 times

pentium75 11 months, 1 week ago

Selected Answer: A

Occasional cost for retrieval from Glacier is nothing compared to the huge storage cost savings compared to C. Still meets the five minute requirement.

upvoted 2 times

master9 11 months, 2 weeks ago

Selected Answer: C

The retrieval price will play an important role here. I selected the "C" option because in "Glacier and use Expedited retrievals" its around $0.004 per GB/month and for STD-IA $0.0125 per GB/month
https://www.cloudforecast.io/blog/aws-s3-pricing-and-optimization-guide/

upvoted 1 times

pentium75 11 months, 1 week ago

But they "will rarely need to restore their files", thus the low cost for occasional expedited retrievals will be nothing compared to the huge storage cost savings.

upvoted 2 times

ngo01214 1 year, 1 month ago

s3 expedited can only be applied on glacier flexible retrieval storage class and s3 intelligent tiering archive access tier. so the answer should be C

upvoted 3 times

pentium75 11 months, 1 week ago

A mentions "G3 Glacier" which has been renamed to "S3 Glacier Flexible Retrieval" and meets the requirements.

upvoted 2 times

Smart 1 year, 3 months ago

Selected Answer: A

I am going with option A, but it is a poorly written question. "For all but the largest archives (more than 250 MB), data accessed by using Expedited retrievals is typically made available within 1–5 minutes. "

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

Answer - A
Fast availability: Although retrieval times for objects stored in Amazon S3 Glacier typically range from minutes to hours, you can use the Expedited retrievals option to expedite access to your archives. By using Expedited retrievals, the files can be made available in a maximum of five minutes when needed. However, Expedited retrievals do incur higher costs compared to standard retrievals.

upvoted 2 times

hsinchang 1 year, 4 months ago

Selected Answer: A

Expedited retrievals are designed for urgent requests and can provide access to data in as little as 1-5 minutes for most archive objects. Standard retrievals typically finish within 3-5 hours for objects stored in the S3 Glacier Flexible Retrieval storage class or S3 Intelligent-Tiering Archive Access tier. These retrievals typically finish within 12 hours for objects stored in the S3 Glacier

Deep Archive storage class or S3 Intelligent-Tiering Deep Archive Access tier. So A.
upvoted 3 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: A

Expedited retrievals allow you to quickly access your data that's stored in the S3 Glacier Flexible Retrieval storage class or the S3 Intelligent-Tiering Archive Access tier when occasional urgent requests for restoring archives are required. Data accessed by using Expedited retrievals is typically made available within 1–5 minutes.
upvoted 2 times

MrAWSAssociate 1 year, 5 months ago

Selected Answer: A

A for sure!
upvoted 2 times

Doyin8807 1 year, 6 months ago

C because A is not the most cost effective
upvoted 1 times

luiscc 1 year, 6 months ago

Selected Answer: A

Expedited retrieval typically takes 1-5 minutes to retrieve data, making it suitable for the company's requirement of having the files available in a maximum of five minutes.
upvoted 5 times

Efren 1 year, 6 months ago

Selected Answer: A

Glacier expedite
upvoted 3 times

EA100 1 year, 6 months ago

Answer - A
Fast availability: Although retrieval times for objects stored in Amazon S3 Glacier typically range from minutes to hours, you can use the Expedited retrievals option to expedite access to your archives. By using Expedited retrievals, the files can be made available in a maximum of five minutes when needed. However, Expedited retrievals do incur higher costs compared to standard retrievals.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #486 *Topic 1*

A company is building a three-tier application on AWS. The presentation tier will serve a static website The logic tier is a containerized application. This application will store data in a relational database. The company wants to simplify deployment and to reduce operational costs.

Which solution will meet these requirements?

A. Use Amazon S3 to host static content. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate for compute power. Use a managed Amazon RDS cluster for the database. **Most Voted**

B. Use Amazon CloudFront to host static content. Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 for compute power. Use a managed Amazon RDS cluster for the database.

C. Use Amazon S3 to host static content. Use Amazon Elastic Kubernetes Service (Amazon EKS) with AWS Fargate for compute power. Use a managed Amazon RDS cluster for the database.

D. Use Amazon EC2 Reserved Instances to host static content. Use Amazon Elastic Kubernetes Service (Amazon EKS) with Amazon EC2 for compute power. Use a managed Amazon RDS cluster for the database.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

Yadav_Sanjay Highly Voted 1 year, 6 months ago

Selected Answer: A

ECS is slightly cheaper than EKS

upvoted 12 times

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: A

Amazon S3 is a highly scalable and cost-effective storage service that can be used to host static website content. It provides durability, high availability, and low latency access to the static files.

Amazon ECS with AWS Fargate eliminates the need to manage the underlying infrastructure. It allows you to run containerized applications without provisioning or managing EC2 instances. This reduces operational overhead and provides scalability.

By using a managed Amazon RDS cluster for the database, you can offload the management tasks such as backups, patching, and monitoring to AWS. This reduces the operational burden and ensures high availability and durability of the database.
upvoted 5 times

awsgeek75 Most Recent 11 months ago

Selected Answer: A

B: CloudFront = Extra cost for something they don't want (CDN)
C: Kubernetes is more operationally complex than ECS containers on Fargate.
D: EC2 expensive
A: S3 is cheap for static content. ECS with Fargate is easiest implantation. Managed RDS is very low op overhead
upvoted 5 times

wsdasdasdqwdaw 1 year, 1 month ago

Why not B ?
upvoted 1 times

wsdasdasdqwdaw 1 year, 1 month ago

Aaa I got it. With CF we are adding additional cost => A.
upvoted 2 times

cyber_bedouin 1 year ago

A is better because ECS Fargate = "containerized application"
upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

Use Amazon S3 to host static content. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate for compute power. Use a managed Amazon RDS cluster for the database.
upvoted 3 times

jaydesai8 1 year, 4 months ago

Selected Answer: A

S3= hosting static contents
Ecs = Little cheaper than EKS
RDS = Database
upvoted 4 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: A

Use Amazon S3 to host static content. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate for compute power. Use a managed Amazon RDS cluster for the database
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #487 *Topic 1*

A company seeks a storage solution for its application. The solution must be highly available and scalable. The solution also must function as a file system be mountable by multiple Linux instances in AWS and on premises through native protocols, and have no minimum size requirements. The company has set up a Site-to-Site VPN for access from its on-premises network to its VPC.

Which storage solution meets these requirements?

A. Amazon FSx Multi-AZ deployments

B. Amazon Elastic Block Store (Amazon EBS) Multi-Attach volumes

C. Amazon Elastic File System (Amazon EFS) with multiple mount targets **Most Voted**

D. Amazon Elastic File System (Amazon EFS) with a single mount target and multiple access points

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

Felix_br Highly Voted 1 year, 6 months ago

Selected Answer: C

The other options are incorrect for the following reasons:

A. Amazon FSx Multi-AZ deployments Amazon FSx is a managed file system service that provides access to file systems that are hosted on Amazon EC2 instances. Amazon FSx does not support native protocols, such as NFS.
B. Amazon Elastic Block Store (Amazon EBS) Multi-Attach volumes Amazon EBS is a block storage service that provides durable, block-level storage volumes for use with Amazon EC2 instances. Amazon EBS Multi-Attach volumes can be attached to multiple EC2 instances at the same time, but they cannot be mounted by multiple Linux instances through native protocols, such as NFS.
D. Amazon Elastic File System (Amazon EFS) with a single mount target and multiple access points A single mount target can only be used to mount the file system on a single EC2 instance. Multiple access points are used to provide access to the file system from different VPCs.

upvoted 12 times

dkw2342 9 months ago

"A single mount target can only be used to mount the file system on a single EC2 instance. Multiple access points are used to

"A single mount target can only be used to mount the file system on a single EC2 instance. Multiple access points are used to provide access to the file system from different VPCs."

This is clearly wrong. You can have exactly one EFS mount target per subnet (AZ), and of course this mount target can be used by many clients (EC2 instances, containers etc.) - see diagram here for example: https://docs.aws.amazon.com/efs/latest/ug/accessing-fs.html

In my opinion, C and D are equally valid answers.

upvoted 2 times

Salilgen 5 months, 2 weeks ago

D is not valid answer because "the solution must be highly available" then it must be multi-AZ: every AZ requires a mount target

upvoted 1 times

unbendable 1 year, 1 month ago

Amazon FSx ONTAP supports clients mounting it with NFS. https://docs.aws.amazon.com/fsx/latest/ONTAPGuide/attach-linux-client.html. Though A is not clear about which FSx product is used

upvoted 2 times

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: C

Amazon EFS is a fully managed file system service that provides scalable, shared storage for Amazon EC2 instances. It supports the Network File System version 4 (NFSv4) protocol, which is a native protocol for Linux-based systems. EFS is designed to be highly available, durable, and scalable.

upvoted 9 times

awsgeek75 Most Recent 11 months ago

Selected Answer: C

A: FSx is a File Server, not a mountable file system
B: EBS can't be mounted on on-prem devices
D: Access points are not same as mount points
C: EFS support multi mount targets and on-prem devices: https://docs.aws.amazon.com/efs/latest/ug/mounting-fs-mount-helper-direct.html

upvoted 3 times

iwannabeawsgod 1 year, 1 month ago

EFS POSIX LINUX

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

C. Amazon Elastic File System (Amazon EFS) with multiple mount targets

upvoted 3 times

boubie44 1 year, 6 months ago

i don't understand why not D?

upvoted 1 times

lucdt4 1 year, 6 months ago

the requirement is mountable by multiple Linux
-> C (multiple mount targets)

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #488 *Topic 1*

A 4-year-old media company is using the AWS Organizations all features feature set to organize its AWS accounts. According to the company's finance team, the billing information on the member accounts must not be accessible to anyone, including the root user of the member accounts.

Which solution will meet these requirements?

A. Add all finance team users to an IAM group. Attach an AWS managed policy named Billing to the group.

B. Attach an identity-based policy to deny access to the billing information to all users, including the root user.

C. Create a service control policy (SCP) to deny access to the billing information. Attach the SCP to the root organizational unit (OU). Most Voted

D. Convert from the Organizations all features feature set to the Organizations consolidated billing feature set.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: C

Service Control Policies (SCP): SCPs are an integral part of AWS Organizations and allow you to set fine-grained permissions on the organizational units (OUs) within your AWS Organization. SCPs provide central control over the maximum permissions that can be granted to member accounts, including the root user.

Denying Access to Billing Information: By creating an SCP and attaching it to the root OU, you can explicitly deny access to billing information for all accounts within the organization. SCPs can be used to restrict access to various AWS services and actions, including billing-related services.

Granular Control: SCPs enable you to define specific permissions and restrictions at the organizational unit level. By denying access to billing information at the root OU, you can ensure that no member accounts, including root users, have access to the billing information.

upvoted 8 times

Kiki_Pass Highly Voted 1 year, 4 months ago

but SCP do not apply to the management account (full admin power)?

upvoted 5 times

TwinSpark 7 months ago

i can understand this information coming from the famous course in udemy. I tought same, but after some research i now think it is a wrong information.
"SCPs affect all users and roles in attached accounts, including the root user. The only exceptions are those described in Tasks and entities not restricted by SCPs."
https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps.html#:~:text=SCPs%20affect%20all%20users%20and,affect%20any%20service%2Dlinked%20role.

upvoted 3 times

potomac Most Recent 1 year, 1 month ago

Selected Answer: C

SCP is for authorization

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

C. Create a service control policy (SCP) to deny access to the billing information. Attach the SCP to the root organizational unit (OU)

upvoted 3 times

PRASAD180 1 year, 5 months ago

C Crt 100%

upvoted 2 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: C

Service control policy are a type of organization policy that you can use to manage permissions in your organization. SCPs offer central control over the maximum available permissions for all accounts in your organization. SCPs help you to ensure your accounts stay within your organization's access control guidelines. SCPs are available only in an organization that has all features enabled.

upvoted 4 times

Abrar2022 1 year, 6 months ago

By denying access to billing information at the root OU, you can ensure that no member accounts, including root users, have access to the billing information.

upvoted 2 times

nosense 1 year, 6 months ago

Selected Answer: C

c for me

upvoted 2 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #489 *Topic 1*

An ecommerce company runs an application in the AWS Cloud that is integrated with an on-premises warehouse solution. The company uses Amazon Simple Notification Service (Amazon SNS) to send order messages to an on-premises HTTPS endpoint so the warehouse application can process the orders. The local data center team has detected that some of the order messages were not received.

A solutions architect needs to retain messages that are not delivered and analyze the messages for up to 14 days.

Which solution will meet these requirements with the LEAST development effort?

A. Configure an Amazon SNS dead letter queue that has an Amazon Kinesis Data Stream target with a retention period of 14 days.

B. Add an Amazon Simple Queue Service (Amazon SQS) queue with a retention period of 14 days between the application and Amazon SNS.

C. Configure an Amazon SNS dead letter queue that has an Amazon Simple Queue Service (Amazon SQS) target with a retention period of 14 days. **Most Voted**

D. Configure an Amazon SNS dead letter queue that has an Amazon DynamoDB target with a TTL attribute set for a retention period of 14 days.

Correct Answer: *C*

Community vote distribution

C (80%) B (20%)

Comments

pentium75 Highly Voted 11 months, 1 week ago

Selected Answer: C

"Configuring an Amazon SNS dead-letter queue for a subscription ...
A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully", this is exactly what C says. https://docs.aws.amazon.com/sns/latest/dg/sns-configure-dead-letter-queue.html

B, an SQS queue "between the application and Amazon SNS" would change the application logic. SQS cannot push messages to the "on-premises https endpoint", rather the destination would have to retrieve messages from the queue. Besides, option B would eventually deliver the messages that failed on the first attempt, which is NOT what is asked for. The goal is to retain undeliverable messages for analysis (NOT to deliver them), and this is typically achieved with a dead letter queue.

upvoted 11 times

awsgeek75 Highly Voted 11 months ago

Selected Answer: C

LEAST development effort!
A: Custom dead letter queue using Kinesis Data Stream (laughable solution!) so lots of coding
B: Change app logic to put SQS between SNS and the app. Also too much coding
D: Same as A, too much code change
C: SNS dead letter queue is by default a SQS que so no coding required

upvoted 5 times

osmk Most Recent 10 months, 4 weeks ago

A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully.https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html

upvoted 3 times

Mikado211 11 months, 3 weeks ago

Selected Answer: C

Problem here SNS dead letter queue is a SQS queue, so technically speaking both B and C are right. But I suppose that they want us to speak about SNS dead letter queue, that nobody do... meh, frustrating.

upvoted 4 times

Mikado211 11 months, 3 weeks ago

Aaaah ok.

So with B == you place the SQS queue between the application and the SNS topic
with C == you place the SQS queue as a DLQ for the SNS topic

Of course it's C !

upvoted 6 times

aws94 12 months ago

Selected Answer: C

https://docs.aws.amazon.com/sns/latest/dg/sns-configure-dead-letter-queue.html

upvoted 3 times

daniel1 1 year, 1 month ago

Selected Answer: C

GPT4 to the rescue:
The most appropriate solution would be to configure an Amazon SNS dead letter queue with an Amazon Simple Queue Service (Amazon SQS) target with a retention period of 14 days (Option C). This setup would ensure that any undelivered messages are retained in the SQS queue for up to 14 days for analysis, with minimal development effort required.

upvoted 2 times

ealpuche 1 year ago

ChatGPT is not a reliable source.

upvoted 9 times

Wayne23Fang 1 year, 1 month ago

Selected Answer: B

I like (B) since it is put SQS before SNS so we could prepare for retention. (C) dead letter Queue is kind of "rescue" effort. Also (C) should mention reprocessing dead letter.

upvoted 1 times

pentium75 11 months, 1 week ago

"Reprocessing dead letters" is not desired here. They want to "retain messages that are not delivered and analyze the messages for up to 14 days", which is what C does.

upvoted 1 times

thanhnv142 1 year, 1 month ago

C is correct. It used a combination of SNS and SQS so it better than B.

upvoted 2 times

iwannabeawsgod 1 year, 1 month ago

Selected Answer: C

C is the answer

upvoted 2 times

Devsin2000 1 year, 2 months ago

B is correct Answer. SQS Retain messages in queues for up to 14 days
C is incorrect because there is nothing called Amazon SNS dead letter queue

upvoted 2 times

pentium75 11 months, 1 week ago

C "Configure an Amazon SNS dead letter queue"
AWS "Configuring an Amazon SNS dead-letter queue"
https://docs.aws.amazon.com/sns/latest/dg/sns-configure-dead-letter-queue.html

upvoted 2 times

RDM10 1 year, 2 months ago

https://docs.aws.amazon.com/sns/latest/dg/sns-configure-dead-letter-queue.html

upvoted 6 times

lemur88 1 year, 3 months ago

Selected Answer: C

https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

C. Configure an Amazon SNS dead letter queue that has an Amazon Simple Queue Service (Amazon SQS) target with a retention period of 14 days.
By using an Amazon SQS queue as the target for the dead letter queue, you ensure that the undelivered messages are reliably stored in a queue for up to 14 days. Amazon SQS allows you to specify a retention period for messages, which meets the retention requirement without additional development effort.

upvoted 2 times

mtmayer 1 year, 3 months ago

Selected Answer: B

Dead Letter is a SQS feature not SNS.
A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully. Messages that can't be delivered due to client errors or server errors are held in the dead-letter queue for further analysis or reprocessing. For more information, see Configuring an Amazon SNS dead-letter queue for a subscription and Amazon SNS message delivery retries.
https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html

upvoted 3 times

pentium75 11 months, 1 week ago

"See Configuring an Amazon SNS (!) dead-letter queue", exactly, thus C.

upvoted 2 times

xyb 1 year, 4 months ago

Selected Answer: B

In SNS, DLQs store the messages that failed to be delivered to subscribed endpoints. For more information, see Amazon SNS Dead-Letter Queues.

In SQS, DLQs store the messages that failed to be processed by your consumer application. This failure mode can happen when producers and consumers fail to interpret aspects of the protocol that they use to communicate. In that case, the consumer receives the message from the queue, but fails to process it, as the message doesn't have the structure or content that the consumer expects. The consumer can't delete the message from the queue either. After exhausting the receive count in the redrive policy, SQS can sideline the message to the DLQ. For more information, see Amazon SQS Dead-Letter Queues.

https://aws.amazon.com/blogs/compute/designing-durable-serverless-apps-with-dlqs-for-amazon-sns-amazon-sqs-aws-lambda/

upvoted 2 times

pentium75 11 months, 1 week ago

"Configuring an Amazon SNS dead-letter queue for a subscription

A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully. "

upvoted 2 times

TariqKipkemei 1 year, 5 months ago

C is best to handle this requirement. Although good to note that dead-letter queue is an SQS queue.

"A dead-letter queue is an Amazon SQS queue that an Amazon SNS subscription can target for messages that can't be delivered to subscribers successfully. Messages that can't be delivered due to client errors or server errors are held in the dead-letter queue for further analysis or reprocessing."

https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html#:~:text=A%20dead%2Dletter%20queue%20is%20an%20Amazon%20SQS%20queue

upvoted 2 times

Felix_br 1 year, 6 months ago

C - Amazon SNS dead letter queues are used to handle messages that are not delivered to their intended recipients. When a message is sent to an Amazon SNS topic, it is first delivered to the topic's subscribers. If a message is not delivered to any of the subscribers, it is sent to the topic's dead letter queue.

Amazon SQS is a fully managed message queuing service that enables you to decouple and scale microservices, distributed systems, and serverless applications. Amazon SQS queues can be configured to have a retention period, which is the amount of time that messages will be kept in the queue before they are deleted.

To meet the requirements of the company, you can configure an Amazon SNS dead letter queue that has an Amazon SQS target with a retention period of 14 days. This will ensure that any messages that are not delivered to the on-premises warehouse application will be stored in the Amazon SQS queue for up to 14 days. The company can then analyze the messages in the Amazon SQS queue to determine why they were not delivered.

upvoted 3 times

Yadav_Sanjay 1 year, 6 months ago

Selected Answer: C

https://docs.aws.amazon.com/sns/latest/dg/sns-dead-letter-queues.html

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #490 *Topic 1*

A gaming company uses Amazon DynamoDB to store user information such as geographic location, player data, and leaderboards. The company needs to configure continuous backups to an Amazon S3 bucket with a minimal amount of coding. The backups must not affect availability of the application and must not affect the read capacity units (RCUs) that are defined for the table.

Which solution meets these requirements?

A. Use an Amazon EMR cluster. Create an Apache Hive job to back up the data to Amazon S3.

B. Export the data directly from DynamoDB to Amazon S3 with continuous backups. Turn on point-in-time recovery for the table. **Most Voted**

C. Configure Amazon DynamoDB Streams. Create an AWS Lambda function to consume the stream and export the data to an Amazon S3 bucket.

D. Create an AWS Lambda function to export the data from the database tables to Amazon S3 on a regular basis. Turn on point-in-time recovery for the table.

Correct Answer: *B*

Community vote distribution

B (89%) C (11%)

Comments

elmogy Highly Voted 1 year, 6 months ago

Selected Answer: B

Continuous backups is a native feature of DynamoDB, it works at any scale without having to manage servers or clusters and allows you to export data across AWS Regions and accounts to any point-in-time in the last 35 days at a per-second granularity. Plus, it doesn't affect the read capacity or the availability of your production tables.

https://aws.amazon.com/blogs/aws/new-export-amazon-dynamodb-table-data-to-data-lake-amazon-s3/

upvoted 12 times

awsgeek75 Highly Voted 11 months ago

Selected Answer: B

A: Impacts RCU

C: Requires coding of Lambda to read from stream to S3
D: More coding in Lambda
B: AWS Managed solution with no coding
upvoted 5 times

potomac Most Recent 1 year, 1 month ago

Selected Answer: B

DynamoDB export to S3 is a fully managed solution for exporting DynamoDB data to an Amazon S3 bucket at scale.
upvoted 4 times

baba365 1 year, 2 months ago

A DynamoDB stream is an ordered flow of information about changes to items in a DynamoDB table… for C.U.D events (Create, Update, Delete) and its logs are retained for only 24hrs .
upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Export the data directly from DynamoDB to Amazon S3 with continuous backups. Turn on point-in-time recovery for the table.
upvoted 3 times

ukivanlamlpi 1 year, 4 months ago

Selected Answer: C

continous backup, no impact to availability ==> DynamoDB stream
B. export is one off, noy continuous and demand on read capacity
upvoted 4 times

hsinchang 1 year, 4 months ago

minimal amount of coding rules out Lambda
upvoted 4 times

Chris22usa 1 year, 5 months ago

ChatGpt answer is C and it indicates continuous backup process uses DynamoDB stream actually
upvoted 2 times

pentium75 11 months, 1 week ago

ChatGPT is usually wrong on these topics.
upvoted 2 times

Gajendr 11 months, 3 weeks ago

Wrong.
"DynamoDB full exports are charged based on the size of the DynamoDB table (table data and local secondary indexes) at the point in time for which the export is done. DynamoDB incremental exports are charged based on the size of data processed from your continuous backups for the time period being exported."
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/S3DataExport.HowItWorks.html
upvoted 1 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: B

Using DynamoDB table export, you can export data from an Amazon DynamoDB table from any time within your point-in-time recovery window to an Amazon S3 bucket. Exporting a table does not consume read capacity on the table, and has no impact on table performance and availability.
upvoted 1 times

norris81 1 year, 6 months ago

Selected Answer: B

https://repost.aws/knowledge-center/back-up-dynamodb-s3
https://aws.amazon.com/blogs/aws/new-amazon-dynamodb-continuous-backups-and-point-in-time-recovery-pitr/
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Streams.Lambda.html

There is no edit
upvoted 2 times

cloudenthusiast 1 year, 6 months ago

Selected Answer: B

Continuous Backups: DynamoDB provides a feature called continuous backups, which automatically backs up your table data. Enabling continuous backups ensures that your table data is continuously backed up without the need for additional coding or manual interventions.

Export to Amazon S3: With continuous backups enabled, DynamoDB can directly export the backups to an Amazon S3 bucket. This eliminates the need for custom coding to export the data.

Minimal Coding: Option B requires the least amount of coding effort as continuous backups and the export to Amazon S3 functionality are built-in features of DynamoDB.

No Impact on Availability and RCUs: Enabling continuous backups and exporting data to Amazon S3 does not affect the availability of your application or the read capacity units (RCUs) defined for the table. These operations happen in the background and do not impact the table's performance or consume additional RCUs.
upvoted 4 times

Efren 1 year, 6 months ago

Selected Answer: B

DynamoDB Export to S3 feature
Using this feature, you can export data from an Amazon DynamoDB table anytime within your point-in-time recovery window to an Amazon S3 bucket.
upvoted 3 times

Efren 1 year, 6 months ago

B also for me
upvoted 3 times

norris81 1 year, 6 months ago

https://repost.aws/knowledge-center/back-up-dynamodb-s3
https://aws.amazon.com/blogs/aws/new-amazon-dynamodb-continuous-backups-and-point-in-time-recovery-pitr/
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Streams.Lambda.html
upvoted 1 times

Efren 1 year, 6 months ago

you could mention what is the best answer from you :)
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #491 *Topic 1*

A solutions architect is designing an asynchronous application to process credit card data validation requests for a bank. The application must be secure and be able to process each request at least once.

Which solution will meet these requirements MOST cost-effectively?

A. Use AWS Lambda event source mapping. Set Amazon Simple Queue Service (Amazon SQS) standard queues as the event source. Use AWS Key Management Service (SSE-KMS) for encryption. Add the kms:Decrypt permission for the Lambda execution role. **Most Voted**

B. Use AWS Lambda event source mapping. Use Amazon Simple Queue Service (Amazon SQS) FIFO queues as the event source. Use SQS managed encryption keys (SSE-SQS) for encryption. Add the encryption key invocation permission for the Lambda function.

C. Use the AWS Lambda event source mapping. Set Amazon Simple Queue Service (Amazon SQS) FIFO queues as the event source. Use AWS KMS keys (SSE-KMS). Add the kms:Decrypt permission for the Lambda execution role.

D. Use the AWS Lambda event source mapping. Set Amazon Simple Queue Service (Amazon SQS) standard queues as the event source. Use AWS KMS keys (SSE-KMS) for encryption. Add the encryption key invocation permission for the Lambda function.

Correct Answer: *A*

Community vote distribution

A (77%) B (21%) C

Comments

elmogy Highly Voted 2 years ago

Selected Answer: A

SQS FIFO is slightly more expensive than standard queue
https://calculator.aws/#/addService/SQS

I would still go with the standard because of the keyword "at least once" because FIFO process "exactly once". That leaves us with A and D, I believe that lambda function only needs to decrypt so I would choose A

upvoted 12 times

pentium75 Highly Voted 1 year, 5 months ago

Selected Answer: A

Selected Answer: A

"Process each request at least once" = Standard queue, rules out B and C which use more expensive FIFO queue

Permissions are added to Lambda execution roles, not Lambda functions, thus D is out.

upvoted 11 times

emakid Most Recent 11 months, 2 weeks ago

Selected Answer: A

Use AWS Lambda event source mapping. Set Amazon Simple Queue Service (Amazon SQS) standard queues as the event source. Use AWS Key Management Service (SSE-KMS) for encryption. Add the kms
permission for the Lambda execution role.

upvoted 2 times

JackyCCK 1 year, 3 months ago

D is not FIFO either

upvoted 1 times

EdenWang 1 year, 6 months ago

Selected Answer: B

With the SSE-SQS encryption type, you do not need to create, manage, or pay for SQS-managed encryption keys.

upvoted 1 times

pentium75 1 year, 5 months ago

And what the hell is "encryption key invocation permission for the Lambda function"?

upvoted 5 times

wsdasdasdqwdaw 1 year, 7 months ago

Initially though it is B, but it is said that the messages should be processed at lest once, not the same order, and Standard SQS is "almost" FIFO, which changed my opinion and I would go with A as correct.

upvoted 4 times

BrijMohan08 1 year, 9 months ago

Selected Answer: A

https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/standard-queues.html

upvoted 5 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

Using SQS FIFO queues ensures each message is processed at least once in order. SSE-SQS provides encryption that is handled entirely by SQS without needing decrypt permissions.

Standard SQS queues (Options A and D) do not guarantee order.

Using KMS keys (Options C and D) requires providing the Lambda role with decrypt permissions, adding complexity.

SQS FIFO queues with SSE-SQS encryption provide orderly, secure, server-side message processing that Lambda can consume without needing to manage decryption. This is the most efficient and cost-effective approach.

upvoted 8 times

Clouddon 1 year, 8 months ago

Amazon SQS offers standard as the default queue type. Standard queues support a nearly unlimited number of API calls per second, per API action (SendMessage, ReceiveMessage, or DeleteMessage). Standard queues support at-least-once message delivery. However, occasionally (because of the highly distributed architecture that allows nearly unlimited throughput), more than one copy of a message might be delivered out of order. Standard queues provide best-effort ordering which ensures that messages are generally delivered in the same order as they're sent.Whereas, FIFO (First-In-First-Out) queues have all the capabilities of the standard queues, but are designed to enhance messaging between applications when the order of operations and events is critical, or where duplicates can't be tolerated. (is correct)

upvoted 5 times

pentium75 1 year, 5 months ago

But permissions are added to Lambda execution roles, not functions
upvoted 5 times

hsinchang 1 year, 10 months ago

Least Privilege Policy leads to A over D.
upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: B

Considering this is credit card validation process, there needs to be a strict 'process exactly once' policy offered by the SQS FIFO, and also SQS already supports server-side encryption with customer-provided encryption keys using the AWS Key Management Service (SSE-KMS) or using SQS-owned encryption keys (SSE-SQS). Both encryption options greatly reduce the operational burden and complexity involved in protecting data. Additionally, with the SSE-SQS encryption type, you do not need to create, manage, or pay for SQS-managed encryption keys.
Therefore option B stands out for me.
upvoted 1 times

TariqKipkemei 1 year, 7 months ago

I retract my answer and change it to A, there is a requirement to process each request 'at least once'. Only standard queues can deliver messages at least once.
There is also a requirement for the most 'cost-effective' option. Standard queues are the cheaper option.

https://aws.amazon.com/sqs/pricing/#:~:text=SQS%20requests%20priced%3F
upvoted 3 times

darren_song 1 year, 11 months ago

Selected Answer: A

https://docs.aws.amazon.com/zh_tw/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-least-privilege-policy.html
upvoted 2 times

Abrar2022 2 years ago

Selected Answer: A

at least once and cost effective suggests SQS standard
upvoted 2 times

Felix_br 2 years ago

Selected Answer: B

Solution B is the most cost-effective solution to meet the requirements of the application.

Amazon Simple Queue Service (SQS) FIFO queues are a good choice for this application because they guarantee that messages are processed in the order in which they are received. This is important for credit card data validation because it ensures that fraudulent transactions are not processed before legitimate transactions.

SQS managed encryption keys (SSE-SQS) are a good choice for encrypting the messages in the SQS queue because they are free to use. AWS Key Management Service (KMS) keys (SSE-KMS) are also a good choice for encrypting the messages, but they do incur a cost.
upvoted 2 times

pentium75 1 year, 5 months ago

"They guarantee that messages are processed in the order in which they are received. This is important" but not asked for!
upvoted 3 times

omoakin 2 years ago

AAAAAAAA
upvoted 1 times

Yadav_Sanjay 2 years ago

Selected Answer: A

should be A. Key word - at least once and cost effective suggests SQS standard

upvoted 3 times

Efren 2 years ago

It has to be default, no FIFO. It doesnt say just one, it says at least once, so that is default queue that is cheaper than FIFO. Between the default options, nto sure to be honest

upvoted 4 times

jayce5 2 years ago

No, when it comes to "credit card data validation," it should be FIFO. If you use the standard approach, there is a chance that people who come after will get processed before those who come first.

upvoted 1 times

pentium75 1 year, 5 months ago

Question clearly says "process each request at least once" which is the description of a standard queue. Your opinion how these transactions should be processed does not matter if it contradicts the requirements given.

Besides, it is about "credit card data validation", NOT payments. Nothing happens if they check twice is your credit card is valid.

upvoted 2 times

awwass 2 years ago

Selected Answer: A

I guess A

upvoted 2 times

awwass 2 years ago

This solution uses standard queues in Amazon SQS, which are less expensive than FIFO queues. It also uses AWS Key Management Service (SSE-KMS) for encryption, which is a cost-effective way to encrypt data at rest and in transit. The kms:Decrypt permission is added to the Lambda execution role to allow it to decrypt messages from the queue

upvoted 2 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #492 *Topic 1*

A company has multiple AWS accounts for development work. Some staff consistently use oversized Amazon EC2 instances, which causes the company to exceed the yearly budget for the development accounts. The company wants to centrally restrict the creation of AWS resources in these accounts.

Which solution will meet these requirements with the LEAST development effort?

A. Develop AWS Systems Manager templates that use an approved EC2 creation process. Use the approved Systems Manager templates to provision EC2 instances.

B. Use AWS Organizations to organize the accounts into organizational units (OUs). Define and attach a service control policy (SCP) to control the usage of EC2 instance types. **Most Voted**

C. Configure an Amazon EventBridge rule that invokes an AWS Lambda function when an EC2 instance is created. Stop disallowed EC2 instance types.

D. Set up AWS Service Catalog products for the staff to create the allowed EC2 instance types. Ensure that staff can deploy EC2 instances only by using the Service Catalog products.

Correct Answer: *B*

Community vote distribution

B (93%) D (7%)

Comments

alexandercamachop Highly Voted 1 year, 6 months ago

Selected Answer: B

Anytime you see Multiple AWS Accounts, and needs to consolidate is AWS Organization. Also anytime we need to restrict anything in an organization, it is SCP policies.

upvoted 8 times

omarshaban Highly Voted 10 months, 3 weeks ago

IN MY EXAM

upvoted 5 times

Cyberkayu Most Recent 11 months, 3 weeks ago

Selected Answer: B

B. Multiple AWS account, consolidate under one AWS Organization, top down policy (SCP) to all member account to restrict EC2 Type.

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Use AWS Organizations to organize the accounts into organizational units (OUs). Define and attach a service control policy (SCP) to control the usage of EC2 instance types.

upvoted 3 times

Ale1973 1 year, 3 months ago

Selected Answer: D

I have a question regarding this answer, what do they mean by "development effort"?:
If they mean the work it takes to implement the solution (using develop as implement), option B achieves the constraint with little administrative overhead (there is less to do to configure this option).
If by "development effort", they mean less effort for the development team, when development team try to deploy instances and gets errors because they are not allowed, this generates overhead. In this case the best option is D.
What did you think?

upvoted 2 times

pentium75 11 months, 1 week ago

"Development effort" = Develop the solution that the question asks for. We don't care about the developers whose permissions we want to restrict.

upvoted 4 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: B

Use AWS Organizations to organize the accounts into organizational units (OUs). Define and attach a service control policy (SCP) to control the usage of EC2 instance types

upvoted 3 times

Blingy 1 year, 6 months ago

BBBBBBBBB

upvoted 2 times

elmogy 1 year, 6 months ago

Selected Answer: B

I would choose B
The other options would require some level of programming or custom resource creation:
A. Developing Systems Manager templates requires development effort
C. Configuring EventBridge rules and Lambda functions requires development effort
D. Creating Service Catalog products requires development effort to define the allowed EC2 configurations.

Option B - Using Organizations service control policies - requires no custom development. It involves:
Organizing accounts into OUs
Creating an SCP that defines allowed/disallowed EC2 instance types
Attaching the SCP to the appropriate OUs
This is a native AWS service with a simple UI for defining and managing policies. No coding or resource creation is needed.
So option B, using Organizations service control policies, will meet the requirements with the least development effort.

upvoted 4 times

cloudenthusiast 1 year, 6 months ago

Selected Answer: B

AWS Organizations: AWS Organizations is a service that helps you centrally manage multiple AWS accounts. It enables you to group accounts into organizational units (OUs) and apply policies across those accounts.

Service Control Policies (SCPs): SCPs in AWS Organizations allow you to define fine-grained permissions and restrictions at the account or OU level. By attaching an SCP to the development accounts, you can control the creation and usage of EC2 instance types.

Least Development Effort: Option B requires minimal development effort as it leverages the built-in features of AWS

Organizations and SCPs. You can define the SCP to restrict the use of oversized EC2 instance types and apply it to the appropriate OUs or accounts.
upvoted 5 times

Efren 1 year, 6 months ago

B for me as well
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #493 *Topic 1*

A company wants to use artificial intelligence (AI) to determine the quality of its customer service calls. The company currently manages calls in four different languages, including English. The company will offer new languages in the future. The company does not have the resources to regularly maintain machine learning (ML) models.

The company needs to create written sentiment analysis reports from the customer service call recordings. The customer service call recording text must be translated into English.

Which combination of steps will meet these requirements? (Choose three.)

A. Use Amazon Comprehend to translate the audio recordings into English.

B. Use Amazon Lex to create the written sentiment analysis reports.

C. Use Amazon Polly to convert the audio recordings into text.

D. Use Amazon Transcribe to convert the audio recordings in any language into text. Most Voted

E. Use Amazon Translate to translate text in any language to English. Most Voted

F. Use Amazon Comprehend to create the sentiment analysis reports. Most Voted

Correct Answer: *DEF*

Community vote distribution

DEF (100%)

Comments

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: DEF

Amazon Transcribe will convert the audio recordings into text, Amazon Translate will translate the text into English, and Amazon Comprehend will perform sentiment analysis on the translated text to generate sentiment analysis reports.

upvoted 7 times

awsgeek75 Highly Voted 11 months ago

Selected Answer: DEF

A: Comprehend cannot translate
B: Lex is like a chatbot so not useful

C: Polly converts text to audio (polly the parrot!) so this is wrong
D: Can convert audio to text
E: Can translate
F: Can do sentiment analysis reports
upvoted 6 times

wsdasdasdqwdaw Most Recent 1 year, 1 month ago

It is: DEF
upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: DEF

D. Use Amazon Transcribe to convert the audio recordings in any language into text.
E. Use Amazon Translate to translate text in any language to English.
F. Use Amazon Comprehend to create the sentiment analysis reports.
upvoted 3 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: DEF

Amazon Transcribe to convert speech to text. Amazon Translate to translate text to english. Amazon Comprehend to perform sentiment analysis on translated text.
upvoted 2 times

HareshPrajapati 1 year, 6 months ago

afree with DEF
upvoted 2 times

Blingy 1 year, 6 months ago

I'd go with DEF too
upvoted 2 times

elmogy 1 year, 6 months ago

Selected Answer: DEF

agree with DEF
upvoted 3 times

Efren 1 year, 6 months ago

agreed as well, weird
upvoted 2 times

Guru4Cloud 1 year, 3 months ago

@efren - It is not weird - This need to know the services for it
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #494 *Topic 1*

A company uses Amazon EC2 instances to host its internal systems. As part of a deployment operation, an administrator tries to use the AWS CLI to terminate an EC2 instance. However, the administrator receives a 403 (Access Denied) error message.

The administrator is using an IAM role that has the following IAM policy attached:

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": ["ec2:TerminateInstances"],
            "Resource": ["*"]
        },
        {
            "Effect": "Deny",
            "Action": ["ec2:TerminateInstances"],
            "Condition": {
                "NotIpAddress": {
                    "aws:SourceIp": [
                        "192.0.2.0/24",
                        "203.0.113.0/24"
                    ]
                }
            },
            "Resource": ["*"]
        }
    ]
}
```

What is the cause of the unsuccessful request?

A. The EC2 instance has a resource-based policy with a Deny statement.

B. The principal has not been specified in the policy statement.

C. The "Action" field does not grant the actions that are required to terminate the EC2 instance.

D. The request to terminate the EC2 instance does not originate from the CIDR blocks 192.0.2.0/24 or 203.0.113.0/24.

Most Voted

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

TariqKipkemei Highly Voted 1 year, 11 months ago

Selected Answer: D

the command is coming from a source IP which is not in the allowed range.

upvoted 5 times

chasingsummer Most Recent 1 year, 5 months ago

Selected Answer: D

I ran a Policy Simulator and indeed, D is right answer.

Here is the JSON policy:

```
{
"Version": "2012-10-17",
"Statement": [
{
"Effect": "Allow",
"Action": "ec2:TerminateInstances",
"Resource": "*"
},
{
"Effect": "Deny",
"Action": "ec2:TerminateInstances",
"Condition": {
"NotIpAddress": {
"aws:SourceIp" : [
"192.0.2.0/24",
"203.0.113.0/24"
]
}
},
"Resource": "*"
}
]
}
```

upvoted 2 times

chasingsummer 1 year, 5 months ago

The condition operator is "NotIpAddress" so I am not sure about D as right answer.

upvoted 2 times

awsgeek75 1 year, 4 months ago

Deny when IP address is not in (NotIPAddress). AWS has a weird way of stating Deny and it almost sound like double negative meaning positive. But read this doc for more clarity:
https://docs.aws.amazon.com/IAM/latest/UserGuide/reference_policies_examples_aws_deny-ip.html

It has the exact same example! Good luck!

upvoted 2 times

LeonSauveterre 6 months ago

This is exactly what I need, thanks!

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

If you want to read more about this, see how it works:
https://docs.aws.amazon.com/IAM/latest/UserGuide/reference_policies_examples_aws_deny-ip.html
Same policy as in this question with almost same use case.
D is correct answer.

upvoted 4 times

elmogy 2 years ago

Selected Answer: D

" aws:SourceIP " indicates the IP address that is trying to perform the action.

upvoted 2 times

nosense 2 years ago

Selected Answer: D

d for sure

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #495 *Topic 1*

A company is conducting an internal audit. The company wants to ensure that the data in an Amazon S3 bucket that is associated with the company's AWS Lake Formation data lake does not contain sensitive customer or employee data. The company wants to discover personally identifiable information (PII) or financial information, including passport numbers and credit card numbers.

Which solution will meet these requirements?

A. Configure AWS Audit Manager on the account. Select the Payment Card Industry Data Security Standards (PCI DSS) for auditing.

B. Configure Amazon S3 Inventory on the S3 bucket Configure Amazon Athena to query the inventory.

C. Configure Amazon Macie to run a data discovery job that uses managed identifiers for the required data types. **Most Voted**

D. Use Amazon S3 Select to run a report across the S3 bucket.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

Blingy Highly Voted 1 year, 6 months ago

Macie = Sensitive PII

upvoted 5 times

elmogy Highly Voted 1 year, 6 months ago

Selected Answer: C

agree with C

upvoted 5 times

awsgeek75 Most Recent 11 months ago

Selected Answer: C

PII or sensitive data = Macie

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

Configure Amazon Macie to run a data discovery job that uses managed identifiers for the required data types.

upvoted 3 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: C

Amazon Macie is a data security service that uses machine learning (ML) and pattern matching to discover and help protect your sensitive data.

upvoted 3 times

cloudenthusiast 1 year, 6 months ago

Selected Answer: C

Amazon Macie is a service that helps discover, classify, and protect sensitive data stored in AWS. It uses machine learning algorithms and managed identifiers to detect various types of sensitive information, including personally identifiable information (PII) and financial information. By configuring Amazon Macie to run a data discovery job with the appropriate managed identifiers for the required data types (such as passport numbers and credit card numbers), the company can identify and classify any sensitive data present in the S3 bucket.

upvoted 5 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #496 *Topic 1*

A company uses on-premises servers to host its applications. The company is running out of storage capacity. The applications use both block storage and NFS storage. The company needs a high-performing solution that supports local caching without re-architecting its existing applications.

Which combination of actions should a solutions architect take to meet these requirements? (Choose two.)

A. Mount Amazon S3 as a file system to the on-premises servers.

B. Deploy an AWS Storage Gateway file gateway to replace NFS storage. **Most Voted**

C. Deploy AWS Snowball Edge to provision NFS mounts to on-premises servers.

D. Deploy an AWS Storage Gateway volume gateway to replace the block storage. **Most Voted**

E. Deploy Amazon Elastic File System (Amazon EFS) volumes and mount them to on-premises servers.

Correct Answer: *BD*

Community vote distribution

BD (100%)

Comments

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: BD

By combining the deployment of an AWS Storage Gateway file gateway and an AWS Storage Gateway volume gateway, the company can address both its block storage and NFS storage needs, while leveraging local caching capabilities for improved performance.

upvoted 7 times

awsgeek75 Most Recent 11 months ago

Selected Answer: BD

A: Not possible
C: Snowball edge is snowball with computing. It's not a NAS!
E: Technically yes but requires VPN or Direct Connect so re-architecture
B & D both use Storage Gateway which can be used as NFS and Block storage
https://aws.amazon.com/storagegateway/

upvoted 4 times

ftaws 11 months, 3 weeks ago

Use the Storage Gateway -> It means that use S3 for storage ?

upvoted 3 times

thanhnv142 1 year, 1 month ago

DE
B is not correct be cause NFS is a file system while storage gw is a storage. To replace a file system, need another file system which is EFS.

upvoted 3 times

Tekk97 1 year ago

That's what i thought. but I think B is work too.

upvoted 2 times

wizcloudifa 7 months, 2 weeks ago

if you focus on the wording of option B its "Storage gateway File gateway" not volume gateway hence it is perfect replacement for NFS files.

upvoted 1 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: BD

Deploy an AWS Storage Gateway file gateway to replace NFS storage
Deploy an AWS Storage Gateway volume gateway to replace the block storage

upvoted 3 times

elmogy 1 year, 6 months ago

Selected Answer: BD

local caching is a key feature of AWS Storage Gateway solution
https://aws.amazon.com/storagegateway/features/
https://aws.amazon.com/blogs/storage/aws-storage-gateway-increases-cache-4x-and-enhances-bandwidth-throttling/#:~:text=AWS%20Storage%20Gateway%20increases%20cache%204x%20and%20enhances,for%20Volume%20Gateway%20customers%20...%205%20Conclusion%20

upvoted 4 times

Piccalo 1 year, 6 months ago

Selected Answer: BD

B and D is the correct answer

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #497 *Topic 1*

A company has a service that reads and writes large amounts of data from an Amazon S3 bucket in the same AWS Region. The service is deployed on Amazon EC2 instances within the private subnet of a VPC. The service communicates with Amazon S3 over a NAT gateway in the public subnet. However, the company wants a solution that will reduce the data output costs.

Which solution will meet these requirements MOST cost-effectively?

A. Provision a dedicated EC2 NAT instance in the public subnet. Configure the route table for the private subnet to use the elastic network interface of this instance as the destination for all S3 traffic.

B. Provision a dedicated EC2 NAT instance in the private subnet. Configure the route table for the public subnet to use the elastic network interface of this instance as the destination for all S3 traffic.

C. Provision a VPC gateway endpoint. Configure the route table for the private subnet to use the gateway endpoint as the route for all S3 traffic. **Most Voted**

D. Provision a second NAT gateway. Configure the route table for the private subnet to use this NAT gateway as the destination for all S3 traffic.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: C

A VPC gateway endpoint allows you to privately access Amazon S3 from within your VPC without using a NAT gateway or NAT instance. By provisioning a VPC gateway endpoint for S3, the service in the private subnet can directly communicate with S3 without incurring data transfer costs for traffic going through a NAT gateway.

upvoted 10 times

awsgeek75 Most Recent 11 months ago

Selected Answer: C

As a rule of thumb, EC2<->S3 in your workload should always try to use a VPC gateway unless there is an explicit restriction (account etc.) which disallows it.

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

Using a VPC endpoint for S3 allows the EC2 instances to access S3 directly over the Amazon network without traversing the internet. This significantly reduces data output charges.

upvoted 3 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: C

use VPC gateway endpoint to route traffic internally and save on costs.

upvoted 2 times

elmogy 1 year, 6 months ago

Selected Answer: C

private subnet needs to communicate with S3 --> VPC endpoint right away

upvoted 3 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #498 *Topic 1*

A company uses Amazon S3 to store high-resolution pictures in an S3 bucket. To minimize application changes, the company stores the pictures as the latest version of an S3 object. The company needs to retain only the two most recent versions of the pictures.

The company wants to reduce costs. The company has identified the S3 bucket as a large expense.

Which solution will reduce the S3 costs with the LEAST operational overhead?

A. Use S3 Lifecycle to delete expired object versions and retain the two most recent versions. **Most Voted**

B. Use an AWS Lambda function to check for older versions and delete all but the two most recent versions.

C. Use S3 Batch Operations to delete noncurrent object versions and retain only the two most recent versions.

D. Deactivate versioning on the S3 bucket and retain the two most recent versions.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: A

S3 Lifecycle policies allow you to define rules that automatically transition or expire objects based on their age or other criteria. By configuring an S3 Lifecycle policy to delete expired object versions and retain only the two most recent versions, you can effectively manage the storage costs while maintaining the desired retention policy. This solution is highly automated and requires minimal operational overhead as the lifecycle management is handled by S3 itself.

upvoted 6 times

VellaDevil Highly Voted 1 year, 5 months ago

Selected Answer: A

A --> "you can also provide a maximum number of noncurrent versions to retain."
https://docs.aws.amazon.com/AmazonS3/latest/userguide/intro-lifecycle-rules.html

upvoted 5 times

awsgeek75 Most Recent 11 months ago

Selected Answer: A

B: Too much work with Lambda
C: Possible but requires lot of work
D: Oxymoron statement... i.e. how do you remove version and retain version at same time without additional overhead? Custom solution may be more work.
A: S3 Lifecycle is designed to retain object and version with set criteria

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

Use S3 Lifecycle to delete expired object versions and retain the two most recent versions.

upvoted 3 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: A

S3 Lifecycle to the rescue...whoooosh

upvoted 3 times

antropaws 1 year, 6 months ago

Selected Answer: A

A is correct.

upvoted 3 times

Konb 1 year, 6 months ago

Selected Answer: A

Agree with LONGMEN

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #499 *Topic 1*

A company needs to minimize the cost of its 1 Gbps AWS Direct Connect connection. The company's average connection utilization is less than 10%. A solutions architect must recommend a solution that will reduce the cost without compromising security.

Which solution will meet these requirements?

A. Set up a new 1 Gbps Direct Connect connection. Share the connection with another AWS account.

B. Set up a new 200 Mbps Direct Connect connection in the AWS Management Console.

C. Contact an AWS Direct Connect Partner to order a 1 Gbps connection. Share the connection with another AWS account.

D. Contact an AWS Direct Connect Partner to order a 200 Mbps hosted connection for an existing AWS account. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (83%) B (17%)

Comments

Abrar2022 Highly Voted 1 year, 6 months ago

Selected Answer: D

Hosted Connection 50 Mbps, 100 Mbps, 200 Mbps,
Dedicated Connection 1 Gbps, 10 Gbps, and 100 Gbps

upvoted 11 times

awsgeek75 Highly Voted 11 months ago

Selected Answer: D

A: Not secure as sharing with another account
B: I don't think this possible as you need ISP to setup Direct Connect
C: Less secure due to sharing
D: Direct connect partners can provide hosted solutions for existing accounts so correct answer

upvoted 5 times

awsgeek75 11 months ago

For B I'm wrong above, it's because you cannot order 200MB connection through management console.

upvoted 2 times

Ravan Most Recent 9 months, 2 weeks ago

Selected Answer: D

No, you cannot directly adjust the speed of an existing Direct Connect connection through the AWS Management Console.

To adjust the speed of an existing Direct Connect connection, you typically need to contact your Direct Connect service provider. They can assist you in modifying the speed of your connection based on your requirements. Depending on the provider, this process may involve submitting a request or contacting their support team to initiate the necessary changes. Keep in mind that adjusting the speed of your Direct Connect connection may also involve contractual and billing considerations.

upvoted 4 times

pentium75 11 months, 1 week ago

Selected Answer: D

< 1 Gbps = Hosted (through partner)

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

If you already have an existing AWS Direct Connect connection configured at 1 Gbps, and you wish to reduce the connection bandwidth to 200 Mbps to minimize costs, you should indeed contact your AWS Direct Connect Partner and request to lower the connection speed to 200 Mbps.

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

I meant D.. DDDDDDDDDD

upvoted 6 times

omoakin 1 year, 6 months ago

BBBBBBBBBBBBB

upvoted 2 times

elmogy 1 year, 6 months ago

Selected Answer: D

company need to setup a cheaper connection (200 M) but B is incorrect because you can only order port speeds of 1, 10, or 100 Gbps
for more flexibility you can go with hosted connection, You can order port speeds between 50 Mbps and 10 Gbps.

https://docs.aws.amazon.com/whitepapers/latest/aws-vpc-connectivity-options/aws-direct-connect.html

upvoted 5 times

cloudenthusiast 1 year, 6 months ago

Selected Answer: B

By opting for a lower capacity 200 Mbps connection instead of the 1 Gbps connection, the company can significantly reduce costs. This solution ensures a dedicated and secure connection while aligning with the company's low utilization, resulting in cost savings.

upvoted 4 times

pentium75 11 months, 1 week ago

But 200M cannot be ordered through Management Console, only partners.

upvoted 3 times

norris81 1 year, 6 months ago

Selected Answer: D

D

For Dedicated Connections, 1 Gbps, 10 Gbps, and 100 Gbps ports are available. For Hosted Connections, connection speeds of 50 Mbps, 100 Mbps, 200 Mbps, 300 Mbps, 400 Mbps, 500 Mbps, 1 Gbps, 2 Gbps, 5 Gbps and 10 Gbps may be ordered from approved AWS Direct Connect Partners. See AWS Direct Connect Partners for more information.

upvoted 5 times

nosense 1 year, 6 months ago

Selected Answer: D

A hosted connection is a lower-cost option that is offered by AWS Direct Connect Partners

upvoted 5 times

Efren 1 year, 6 months ago

Also, there are not 200 MBps direct connection speed.

upvoted 1 times

nosense 1 year, 6 months ago

Hosted Connection 50 Mbps, 100 Mbps, 200 Mbps,
Dedicated Connection 1 Gbps, 10 Gbps, and 100 Gbps
B would require the company to purchase additional hardware or software

upvoted 3 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #500 *Topic 1*

A company has multiple Windows file servers on premises. The company wants to migrate and consolidate its files into an Amazon FSx for Windows File Server file system. File permissions must be preserved to ensure that access rights do not change.

Which solutions will meet these requirements? (Choose two.)

A. Deploy AWS DataSync agents on premises. Schedule DataSync tasks to transfer the data to the FSx for Windows File Server file system. **Most Voted**

B. Copy the shares on each file server into Amazon S3 buckets by using the AWS CLI. Schedule AWS DataSync tasks to transfer the data to the FSx for Windows File Server file system.

C. Remove the drives from each file server. Ship the drives to AWS for import into Amazon S3. Schedule AWS DataSync tasks to transfer the data to the FSx for Windows File Server file system.

D. Order an AWS Snowcone device. Connect the device to the on-premises network. Launch AWS DataSync agents on the device. Schedule DataSync tasks to transfer the data to the FSx for Windows File Server file system. **Most Voted**

E. Order an AWS Snowball Edge Storage Optimized device. Connect the device to the on-premises network. Copy data to the device by using the AWS CLI. Ship the device back to AWS for import into Amazon S3. Schedule AWS DataSync tasks to transfer the data to the FSx for Windows File Server file system.

Correct Answer: *AD*

Community vote distribution

AD (96%) BD

Comments

pentium75 Highly Voted 11 months, 1 week ago

Selected Answer: AD

B, C and E would copy the files to S3 first where permissions would be lost

upvoted 10 times

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: AD

A This option involves deploying DataSync agents on your on-premises file servers and using DataSync to transfer the data

A This option involves deploying DataSync agents on your on-premises file servers and using DataSync to transfer the data directly to the FSx for Windows File Server. DataSync ensures that file permissions are preserved during the migration process.
D
This option involves using an AWS Snowcone device, a portable data transfer device. You would connect the Snowcone device to your on-premises network, launch DataSync agents on the device, and schedule DataSync tasks to transfer the data to FSx for Windows File Server. DataSync handles the migration process while preserving file permissions.
upvoted 8 times

Guru4Cloud Most Recent 1 year, 3 months ago
Selected Answer: BD

Why not - BD?
upvoted 1 times

pentium75 11 months, 1 week ago

B copies the data to S3 first where file permissions would be lost.
upvoted 3 times

Guru4Cloud 1 year, 3 months ago

° This option uses S3 as an intermediary, ensuring that file permissions are preserved during the initial data copy. DataSync can then transfer the data from S3 to FSx while maintaining the permissions.
° This option uses a Snowcone device with DataSync agents to replicate the on-premises permission structure directly to FSx. This approach is suitable for maintaining file permissions during migration.
upvoted 2 times

elmogy 1 year, 6 months ago
Selected Answer: AD

the key is file permissions are preserved during the migration process. only datasync supports that
upvoted 5 times

coolkidsclubvip 1 year, 3 months ago

Bro,all 5 answers mentioned Datasync.....
upvoted 5 times

Devsin2000 1 year, 2 months ago

Yes but AD have only DataSync, whereas others have others have AWS CLI used.
upvoted 3 times

nosense 1 year, 6 months ago
Selected Answer: AD

Option B would require copy the data to Amazon S3 before transferring it to Amazon FSx for Windows File Server
Option C would require the company to remove the drives from each file server and ship them to AWS
upvoted 2 times

barracouto 1 year, 3 months ago

Also, S3 doesn't retain permissions because it isn't a file system.
upvoted 4 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #501 *Topic 1*

A company wants to ingest customer payment data into the company's data lake in Amazon S3. The company receives payment data every minute on average. The company wants to analyze the payment data in real time. Then the company wants to ingest the data into the data lake.

Which solution will meet these requirements with the MOST operational efficiency?

A. Use Amazon Kinesis Data Streams to ingest data. Use AWS Lambda to analyze the data in real time.

B. Use AWS Glue to ingest data. Use Amazon Kinesis Data Analytics to analyze the data in real time.

C. Use Amazon Kinesis Data Firehose to ingest data. Use Amazon Kinesis Data Analytics to analyze the data in real time. **Most Voted**

D. Use Amazon API Gateway to ingest data. Use AWS Lambda to analyze the data in real time.

Correct Answer: *C*

Community vote distribution

C (95%) A (5%)

Comments

Axeashes Highly Voted 1 year, 5 months ago

Kinesis Data Firehose is near real time (min. 60 sec). - The question is focusing on real time processing/analysis + efficiency -> Kinesis Data Stream is real time ingestion.
https://www.amazonaws.cn/en/kinesis/data-firehose/#:~:text=Near%20real%2Dtime,is%20sent%20to%20the%20service.
upvoted 12 times

Axeashes 1 year, 5 months ago

Unless the intention is real time analytics not real time ingestion !
upvoted 5 times

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: C

By leveraging the combination of Amazon Kinesis Data Firehose and Amazon Kinesis Data Analytics, you can efficiently ingest and analyze the payment data in real time without the need for manual processing or additional infrastructure management. This solution provides a streamlined and scalable approach to handle continuous data ingestion and analysis requirements.

upvoted 11 times

zdi561 Most Recent 4 months, 2 weeks ago

Selected Answer: C

The question says that the data is analyzed first then ingest into s3 for data lake

upvoted 1 times

wizcloudifa 7 months, 2 weeks ago

Selected Answer: C

Kinesis Firehouse = ingesting
Kinesis Datastreams = storing
Kinesis analytics = doing analysis

upvoted 6 times

awsgeek75 11 months ago

Selected Answer: C

Data is stored on S3 so real-time data analytics can be done with Kinesis Data Analytics which rules out Lambda solutions (A and D) as they are more operationally complex.
B is not useful it is more of ETL.

Firehose is actually to distribute data but given that company is already receiving data somehow so Firehose can basically distribute it to S3 with minimum latency. I have to admit this was confusing. I would have used Kinesis Streams to store on S3 and Data Analytics but combination is confusing!

upvoted 3 times

Mr_Marcus 6 months, 1 week ago

"Data is stored on S3..."
Nope. Re-read the first sentence. S3 is the destination, not the source.
The task is to ingest, analyze in real time, and store in S3.

upvoted 2 times

1rob 1 year ago

Selected Answer: C

"payment data every minute on average" is a good-to-go- for firehose.
Also firehose is more operational efficient compared to Data Streams.

upvoted 3 times

lucasbg 1 year ago

Selected Answer: A

I think this is A. The purpose of Firehose is to ingest and deliver to a data store, no to an analytics service. And in fact you can use lambda for real time analysis, such I find A more aligned.

upvoted 2 times

pentium75 11 months, 1 week ago

But developing and maintaining a custom Lambda function "to analyze the data in real time" is surely not as 'operationally efficient' as using an existing service such as Kinesis Data Analytics.

upvoted 4 times

DDongi 1 year, 1 month ago

Firehose has a 60 sec delay so real time analytics should be without real time data isn't that problematic? Why would you have then real time analytics then in the first place?

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

Kinesis Data Streams focuses on ingesting and storing data streams while Kinesis Data Firehose focuses on delivering data streams to select destinations, as the motive of the question is to do analytics, the answer should be C.

upvoted 4 times

hsinchang 1 year, 4 months ago

[illegible]

Selected Answer: C

Kinesis Data Streams focuses on ingesting and storing data streams while Kinesis Data Firehose focuses on delivering data streams to select destinations, as the motive of the question is to do analytics, the answer should be C.

upvoted 2 times

james2033 1 year, 4 months ago

Selected Answer: C

Quote "Connect with 30+ fully integrated AWS services and streaming destinations such as Amazon Simple Storage Service (S3)" at https://aws.amazon.com/kinesis/data-firehose/ . Amazon Kinesis Data Analystics https://aws.amazon.com/kinesis/data-analytics/

upvoted 2 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: C

Use Kinesis Firehose to capture and deliver the data to Kinesis Analytics to perform analytics.

upvoted 2 times

Anmol_1010 1 year, 6 months ago

Did anyome took tge exam recently,
How many questiona were there

upvoted 2 times

omoakin 1 year, 6 months ago

Can we understand why admin's answers are mostly wrong? Or is this done on purpose?

upvoted 3 times

nosense 1 year, 6 months ago

Selected Answer: C

Amazon Kinesis Data Firehose the most optimal variant

upvoted 4 times

kailu 1 year, 6 months ago

Shouldn't C be more appropriate?

upvoted 5 times

MostofMichelle 1 year, 6 months ago

You're right. I believe the answers are wrong on purpose, so good thing votes can be made on answers and discussions are allowed.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #502 *Topic 1*

A company runs a website that uses a content management system (CMS) on Amazon EC2. The CMS runs on a single EC2 instance and uses an Amazon Aurora MySQL Multi-AZ DB instance for the data tier. Website images are stored on an Amazon Elastic Block Store (Amazon EBS) volume that is mounted inside the EC2 instance.

Which combination of actions should a solutions architect take to improve the performance and resilience of the website? (Choose two.)

A. Move the website images into an Amazon S3 bucket that is mounted on every EC2 instance

B. Share the website images by using an NFS share from the primary EC2 instance. Mount this share on the other EC2 instances.

C. Move the website images onto an Amazon Elastic File System (Amazon EFS) file system that is mounted on every EC2 instance. **Most Voted**

D. Create an Amazon Machine Image (AMI) from the existing EC2 instance. Use the AMI to provision new instances behind an Application Load Balancer as part of an Auto Scaling group. Configure the Auto Scaling group to maintain a minimum of two instances. Configure an accelerator in AWS Global Accelerator for the website

E. Create an Amazon Machine Image (AMI) from the existing EC2 instance. Use the AMI to provision new instances behind an Application Load Balancer as part of an Auto Scaling group. Configure the Auto Scaling group to maintain a minimum of two instances. Configure an Amazon CloudFront distribution for the website. **Most Voted**

Correct Answer: *CE*

Community vote distribution

CE (75%) AE (25%)

Comments

cloudenthusiast Highly Voted 2 years ago

Selected Answer: CE

By combining the use of Amazon EFS for shared file storage and Amazon CloudFront for content delivery, you can achieve improved performance and resilience for the website.

upvoted 14 times

wizcloudifa Highly Voted 1 year, 1 month ago

Selected Answer: CE

First of all you should understand, a website using CMS is a dynamic one not static, so A is out, B is more complicated than C, so C, and between global accelerator and cloudfront, Cloudfront suits better as there is no legacy protocols data(UDP, etc) that needs to be accessed, hence E

upvoted 5 times

foha2012 Most Recent 1 year, 5 months ago

I choose AE. Although I don't know if s3 can be mounted on ec2 ?? Maybe wrong wording. Efs is a better choice but its not a natural selection for strong images.

upvoted 2 times

awsgeek75 1 year, 5 months ago

I made the same mistake but mounting S3 on EC2 is a painful operation so EFS makes more sense (C).
Option E takes care of caching static images on CDN so that problem is solved along with resilience etc.

upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: CE

Not A because you can't mount an S3 bucket on an EC2 instance. You could use a file gateway and share an S3 bucket via NFS and mount that on EC2, but that is not mentioned here and would also not make sense.

upvoted 4 times

seldiora 8 months, 1 week ago

it is possible to mount s3 to EC2, just quite difficult: https://aws.amazon.com/blogs/storage/mounting-amazon-s3-to-an-amazon-ec2-instance-using-a-private-connection-to-s3-file-gateway/

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: CE

You can mount EFS file systems to multiple Amazon EC2 instances remotely and securely without having to log in to the instances by using the AWS Systems Manager Run Command.

upvoted 3 times

wsdasdasdqwdaw 1 year, 7 months ago

A is out of the game for sure. Mount S3 to EC2 ... madness. The question is CE or DE, but it is CE because of AWS Global Accelerator is match with NLB, not ALB as it is staded in option D, thus CE as many of all here noted.

upvoted 5 times

thanhnv142 1 year, 7 months ago

A and E is correct. We have a cloud fromt + S3 combo

upvoted 1 times

wsdasdasdqwdaw 1 year, 7 months ago

S3 can't be mounted on EC2 it is not A for sure!

upvoted 4 times

NickGordon 1 year, 7 months ago

https://aws.amazon.com/blogs/storage/mounting-amazon-s3-to-an-amazon-ec2-instance-using-a-private-connection-to-s3-file-gateway/

upvoted 3 times

pentium75 1 year, 5 months ago

"Using a ... S3 file gateway"

upvoted 2 times

thanhnv142 1 year, 7 months ago

A and E.
C is not correct because You dont mount a new EFS onto existing EC2. If you do that, you have to migrate all exising data in EBS into EFS. Then remove all the EBS. Should never do this.
upvoted 1 times

pentium75 1 year, 5 months ago

I can't follow. EFS provides NFS mount points, how can you not mount those onto existing EC2?
upvoted 2 times

franbarberan 1 year, 8 months ago

Selected Answer: CE

https://bluexp.netapp.com/blog/ebs-efs-amazons3-best-cloud-storage-system
upvoted 2 times

Smart 1 year, 9 months ago

Selected Answer: CE

Not A - S3 cannot be mounted (up until few months ago). Exam does not test for the updates in last 6 months.
upvoted 4 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: AE

You have summarized the reasons why options A and E are the best choices very well.

Migrating static website assets like images to Amazon S3 enables high scalability, durability and shared access across instances. This improves performance.

Using Auto Scaling with load balancing provides elasticity and resilience. Adding a CloudFront distribution further boosts performance through caching and content delivery.
upvoted 2 times

pentium75 1 year, 5 months ago

You can't directly mount an S3 bucket on EC2.
upvoted 2 times

Ale1973 1 year, 10 months ago

Selected Answer: AE

Both options AE and CE would work, but I choose AE, because, on my opinion, S3 is best suited for performance and resilience.
upvoted 3 times

pentium75 1 year, 5 months ago

You can't directly mount an S3 bucket on EC2
upvoted 2 times

MicketyMouse 1 year, 10 months ago

Selected Answer: CE

EFS, unlike EBS, can be mounted across multiple EC2 instances and hence C over A.
upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: AE

Technically both options AE and CE would work. But S3 is best suited for unstructured data, and the key benefit of mounting S3 on EC2 is that it provides a cost-effective alternative of using object storage for applications dealing with large files, as compared to expensive file or block storage. At the same time it provides more performant, scalable and highly available storage for these applications.

Even though there is no mention of 'cost efficient' in this question, in the real world cost is the no.1 factor.
In the exam I believe both options would be a pass.

https://aws.amazon.com/blogs/storage/mounting-amazon-s3-to-an-amazon-ec2-instance-using-a-private-connection-to-s3-

file-gateway/
upvoted 4 times

pentium75 1 year, 5 months ago

You can't directly mount an S3 bucket on EC2, only through file gateway
upvoted 1 times

AshutoshSingh1923 1 year, 11 months ago

Selected Answer: CE

Option C provides moving the website images onto an Amazon EFS file system that is mounted on every EC2 instance. Amazon EFS provides a scalable and fully managed file storage solution that can be accessed concurrently from multiple EC2 instances. This ensures that the website images can be accessed efficiently and consistently by all instances, improving performance
In Option E The Auto Scaling group maintains a minimum of two instances, ensuring resilience by automatically replacing any unhealthy instances. Additionally, configuring an Amazon CloudFront distribution for the website further improves performance by caching content at edge locations closer to the end-users, reducing latency and improving content delivery.
Hence combining these actions, the website's performance is improved through efficient image storage and content delivery
upvoted 3 times

Vadbro7 1 year, 11 months ago

Which answe is correct?the most voted ones or the Suggested answers?
upvoted 1 times

9be0170 11 months ago

the most voted always
upvoted 2 times

mattcl 1 year, 11 months ago

A and E: S3 is perfect for images. Besides is the perfect partner of cloudfront
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #503 *Topic 1*

A company runs an infrastructure monitoring service. The company is building a new feature that will enable the service to monitor data in customer AWS accounts. The new feature will call AWS APIs in customer accounts to describe Amazon EC2 instances and read Amazon CloudWatch metrics.

What should the company do to obtain access to customer accounts in the MOST secure way?

A. Ensure that the customers create an IAM role in their account with read-only EC2 and CloudWatch permissions and a trust policy to the company's account. **Most Voted**

B. Create a serverless API that implements a token vending machine to provide temporary AWS credentials for a role with read-only EC2 and CloudWatch permissions.

C. Ensure that the customers create an IAM user in their account with read-only EC2 and CloudWatch permissions. Encrypt and store customer access and secret keys in a secrets management system.

D. Ensure that the customers create an Amazon Cognito user in their account to use an IAM role with read-only EC2 and CloudWatch permissions. Encrypt and store the Amazon Cognito user and password in a secrets management system.

Correct Answer: *A*

Community vote distribution

A (94%) D (6%)

Comments

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: A

By having customers create an IAM role with the necessary permissions in their own accounts, the company can use AWS Identity and Access Management (IAM) to establish cross-account access. The trust policy allows the company's AWS account to assume the customer's IAM role temporarily, granting access to the specified resources (EC2 instances and CloudWatch metrics) within the customer's account. This approach follows the principle of least privilege, as the company only requests the necessary permissions and does not require long-term access keys or user credentials from the customers.

upvoted 15 times

Piccalo Highly Voted 1 year, 6 months ago

Selected Answer: A

A. Roles give temporary credentials

upvoted 9 times

Efren 1 year, 6 months ago

Agreed . Role is the keyword
upvoted 3 times

awsgeek75 Most Recent 10 months, 3 weeks ago

Selected Answer: A

B: Sharing credentials, even temporary, is insecure
C: Access and secret keys. That won't work and sharing secrets outside of account is not secure for this use case

A: Keyword "trust policy"
D: Again, sharing username and pwd and sharing in any way is not secure
upvoted 2 times

pentium75 11 months, 1 week ago

Selected Answer: A

Not B (would be about access to the company's account, not the customers' accounts)
Not C (storing credentials in a custom system is a big nono)
Not D (Cognito has nothing to do here and "user and password" is terrible)
upvoted 3 times

1rob 1 year ago

Selected Answer: D

The company's infrastructure monitoring service needs to call aws API's in the MOST secure way. So you have to focus on restricting access to the APIs and there is where cognito comes in to play.
upvoted 2 times

awsgeek75 11 months ago

Are you suggesting to restrict CloudWatch API with Cognito roles?
upvoted 2 times

pentium75 11 months, 1 week ago

What is unsecure with A?
upvoted 2 times

1rob 11 months ago

The company runs an infrastructure monitoring service. Nowhere is stated that this service lives in an aws account. So A and C I wouldn't choose.
B is a bit too vague. So I end up with D.
upvoted 1 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

A is the most secure approach for accessing customer accounts.

Having customers create a cross-account IAM role with the appropriate permissions, and configuring the trust policy to allow the monitoring service principal account access, implements secure delegation and least privilege access.
upvoted 2 times