

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #504

Topic 1

A company needs to connect several VPCs in the us-east-1 Region that span hundreds of AWS accounts. The company's networking team has its own AWS account to manage the cloud network.

What is the MOST operationally efficient solution to connect the VPCs?

- A. Set up VPC peering connections between each VPC. Update each associated subnet's route table
- B. Configure a NAT gateway and an internet gateway in each VPC to connect each VPC through the internet
- C. Create an AWS Transit Gateway in the networking team's AWS account. Configure static routes from each VPC.
- D. Deploy VPN gateways in each VPC. Create a transit VPC in the networking team's AWS account to connect to each VPC.

Most Voted

Correct Answer: C

Community vote distribution

C (100%)

Comments

hsinchang Highly Voted 1 year, 4 months ago

Selected Answer: C

The main difference between AWS Transit Gateway and VPC peering is that AWS Transit Gateway is designed to connect multiple VPCs together in a hub-and-spoke model, while VPC peering is designed to connect two VPCs together in a peer-to-peer model.

As we have several VPCs here, the answer should be C.

upvoted 20 times

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: C

AWS Transit Gateway is a highly scalable and centralized hub for connecting multiple VPCs, on-premises networks, and remote networks. It simplifies network connectivity by providing a single entry point and reducing the number of connections required. In this scenario, deploying an AWS Transit Gateway in the networking team's AWS account allows for efficient management and control over the network connectivity across multiple VPCs.

upvoted 7 times

awsgeek75 Most Recent 10 months, 3 weeks ago

Selected Answer: C

A: This option is suggesting hundreds of peering connection for EACH VPC. Nope!
B: NAT gateway is for network translation not VPC interconnectivity so this is wrong
C: Transit GW + static routes will connect all VPCs <https://aws.amazon.com/transit-gateway/>
D: VPN gateway is for on-prem to VPN for a VPC. There is no on-prem here so this is wrong
upvoted 3 times

TariqKipkemei 1 year, 1 month ago

Selected Answer: C

Connect, Monitor and Manage Multiple VPCs in one place = AWS Transit Gateway
upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

C is the most operationally efficient solution for connecting a large number of VPCs across accounts.

Using AWS Transit Gateway allows all the VPCs to connect to a central hub without needing to create a mesh of VPC peering connections between each VPC pair.

This significantly reduces the operational overhead of managing the network topology as new VPCs are added or changed.

The networking team can centrally manage the Transit Gateway routing and share it across accounts using Resource Access Manager.
upvoted 3 times

MirKhobaeb 1 year, 6 months ago

Answer is C
upvoted 2 times

MirKhobaeb 1 year, 6 months ago

A transit gateway is a network transit hub that you can use to interconnect your virtual private clouds (VPCs) and on-premises networks. As your cloud infrastructure expands globally, inter-Region peering connects transit gateways together using the AWS Global Infrastructure. Your data is automatically encrypted and never travels over the public internet.
upvoted 3 times

nosense 1 year, 6 months ago

Selected Answer: C

I voted for c
upvoted 3 times

nosense 1 year, 6 months ago

An AWS Transit Gateway is a highly scalable and secure way to connect VPCs in multiple AWS accounts. It is a central hub that routes traffic between VPCs, on-premises networks, and AWS services.
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #505

Topic 1

A company has Amazon EC2 instances that run nightly batch jobs to process data. The EC2 instances run in an Auto Scaling group that uses On-Demand billing. If a job fails on one instance, another instance will reprocess the job. The batch jobs run between 12:00 AM and 06:00 AM local time every day.

Which solution will provide EC2 instances to meet these requirements MOST cost-effectively?

- A. Purchase a 1-year Savings Plan for Amazon EC2 that covers the instance family of the Auto Scaling group that the batch job uses.
- B. Purchase a 1-year Reserved Instance for the specific instance type and operating system of the instances in the Auto Scaling group that the batch job uses.
- C. Create a new launch template for the Auto Scaling group. Set the instances to Spot Instances. Set a policy to scale out based on CPU usage. **Most Voted**
- D. Create a new launch template for the Auto Scaling group. Increase the instance size. Set a policy to scale out based on CPU usage.

Correct Answer: C

Community vote distribution

C (100%)

Comments

cloudenthusiast **Highly Voted** 1 year, 6 months ago

Selected Answer: C

Purchasing a 1-year Savings Plan (option A) or a 1-year Reserved Instance (option B) may provide cost savings, but they are more suitable for long-running, steady-state workloads. Since your batch jobs run for a specific period each day, using Spot Instances with the ability to scale out based on CPU usage is a more cost-effective choice.

upvoted 14 times

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: C

C is the most cost-effective solution in this scenario.

Using Spot Instances allows EC2 capacity to be purchased at significant discounts compared to On-Demand prices. The auto

scaling group can scale out to add Spot Instances when needed for the batch jobs.

If Spot Instances become unavailable, regular On-Demand Instances will be launched instead to maintain capacity. The potential for interruptions is acceptable since failed jobs can be re-run.

upvoted 7 times

sandordini Most Recent 7 months, 3 weeks ago

Selected Answer: C

Stateless, most cost-effective >> Spot

upvoted 3 times

awsgeek75 11 months ago

Selected Answer: C

You don't need any scaling really as the job runs on another EC2 instance if it fails on first one. A. B. D are all more expensive than C due to spot instance being cheaper than reserved instances.

upvoted 3 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: C

Spot Instances to the rescue....whoosh

upvoted 2 times

wRhIH 1 year, 5 months ago

" If a job fails on one instance, another instance will reprocess the job". This ensures Spot Instances are enough for this case

upvoted 4 times

Abrar2022 1 year, 6 months ago

Selected Answer: C

Since your batch jobs run for a specific period each day, using Spot Instances with the ability to scale out based on CPU usage is a more cost-effective choice.

upvoted 2 times

Blingy 1 year, 6 months ago

C FOR ME COS OF SPOT INSTANCES

upvoted 3 times

udo2020 1 year, 6 months ago

First I think it is B but because of cost saving I think it should be C spot instances.

upvoted 2 times

nosense 1 year, 6 months ago

Selected Answer: C

c for me

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #506

Topic 1

A social media company is building a feature for its website. The feature will give users the ability to upload photos. The company expects significant increases in demand during large events and must ensure that the website can handle the upload traffic from users.


Which solution meets these requirements with the MOST scalability?

- A. Upload files from the user's browser to the application servers. Transfer the files to an Amazon S3 bucket.
- B. Provision an AWS Storage Gateway file gateway. Upload files directly from the user's browser to the file gateway.
- C. Generate Amazon S3 presigned URLs in the application. Upload files directly from the user's browser into an S3 bucket.
- D. Provision an Amazon Elastic File System (Amazon EFS) file system. Upload files directly from the user's browser to the file system.

Most Voted

Correct Answer: C

Community vote distribution



Comments

cloudenthusiast Highly Voted 1 year, 6 months ago

Selected Answer: C

This approach allows users to upload files directly to S3 without passing through the application servers, reducing the load on the application and improving scalability. It leverages the client-side capabilities to handle the file uploads and offloads the processing to S3.

upvoted 16 times

Guru4Cloud Highly Voted 1 year, 3 months ago

Selected Answer: C

C is the best solution to meet the scalability requirements.

Generating S3 presigned URLs allows users to upload directly to S3 instead of application servers. This removes the application servers as a bottleneck for upload traffic.

S3 can scale to handle very high volumes of uploads with no limits on storage or throughput. Using presigned URLs leverages

this scalability.

upvoted 5 times

hro Most Recent 8 months, 2 weeks ago

C - You may use presigned URLs to allow someone to upload an object to your Amazon S3 bucket. Using a presigned URL will allow an upload without requiring another party to have AWS security credentials or permissions.

upvoted 3 times

awsgeek75 11 months ago

Selected Answer: C

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/using-presigned-url.html>

"You can also use presigned URLs to allow someone to upload a specific object to your Amazon S3 bucket. This allows an upload without requiring another party to have AWS security credentials or permissions. "

upvoted 2 times

Goutham4981 1 year ago

Selected Answer: A

S3 presigned url is used for sharing objects from an s3 bucket and not for uploading to an s3 bucket

upvoted 3 times

Murtadhaceit 12 months ago

No. It allows to download and upload.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/using-presigned-url.html>

upvoted 4 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: C

You may use presigned URLs to allow someone to upload an object to your Amazon S3 bucket. Using a presigned URL will allow an upload without requiring another party to have AWS security credentials or permissions.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/PresignedUrlUploadObject.html>

upvoted 2 times

baba365 1 year, 5 months ago

Hello Moderator. This question and answer should be rephrased because:

1. S3 pre-signed URLs are used to share objects FROM S3 buckets
2. How scalable are pre-signed URLs when they are time constrained?

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/ShareObjectPreSignedURL.html>

upvoted 2 times

pentium75 11 months, 1 week ago

Both is wrong

Presigned URLs can be used for upload

The solution is scalable because you can issue thousands of pre-signed URLs, and thousands of users can upload images at the same time.

User wants to upload picture -> server generates presigned URL and sends it to the app -> app uploads file

upvoted 3 times

nosense 1 year, 6 months ago

Selected Answer: C

the most scalable because it allows users to upload files directly to Amazon S3,

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #507


Topic 1

A company has a web application for travel ticketing. The application is based on a database that runs in a single data center in North America. The company wants to expand the application to serve a global user base. The company needs to deploy the application to multiple AWS Regions. Average latency must be less than 1 second on updates to the reservation database.

The company wants to have separate deployments of its web platform across multiple Regions. However, the company must maintain a single primary reservation database that is globally consistent.

Which solution should a solutions architect recommend to meet these requirements?

- A. Convert the application to use Amazon DynamoDB. Use a global table for the center reservation table. Use the correct Regional endpoint in each Regional deployment. **Most Voted**
- B. Migrate the database to an Amazon Aurora MySQL database. Deploy Aurora Read Replicas in each Region. Use the correct Regional endpoint in each Regional deployment for access to the database.
- C. Migrate the database to an Amazon RDS for MySQL database. Deploy MySQL read replicas in each Region. Use the correct Regional endpoint in each Regional deployment for access to the database.
- D. Migrate the application to an Amazon Aurora Serverless database. Deploy instances of the database to each Region. Use the correct Regional endpoint in each Regional deployment to access the database. Use AWS Lambda functions to process event streams in each Region to synchronize the databases.

Correct Answer: A*Community vote distribution***Comments****cloudenthusiast** **Highly Voted** 2 years ago**Selected Answer: A**

Using DynamoDB's global tables feature, you can achieve a globally consistent reservation database with low latency on updates, making it suitable for serving a global user base. The automatic replication provided by DynamoDB eliminates the need for manual synchronization between Regions.

upvoted 21 times

FlvinnHawk 5 months. 4 weeks ago

khatingarun 5 months, 1 week ago

DynamoDB Global Tables rely on eventual consistency for replication between Regions. This means that updates in one Region may take milliseconds to seconds to propagate to other Regions.
<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ReadConsistency.html>

The requirement specifies global consistency, which DynamoDB cannot guarantee.
upvoted 1 times

upliftinghut **Highly Voted** 1 year, 4 months ago

Selected Answer: B

Aurora: less than 1 second: <https://aws.amazon.com/rds/aurora/global-database/>
DynamoDB: from 0.5 to 2.5 second:
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/V2globaltables_HowItWorks.html
upvoted 6 times

TheLaPlanta 1 year, 2 months ago

B doesn't say Aurora Global
upvoted 11 times

MatAlves 8 months, 3 weeks ago

DynamoDB doesn't meet the <1s req though.
upvoted 2 times

tch **Most Recent** 3 months ago

Selected Answer: A

answer should be A.
we can consider if B change to "Amazon Aurora Global Database" ...
upvoted 1 times

tch 3 months ago

Global tables - multi-Region replication for DynamoDB
<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/GlobalTables.htm> Amazon DynamoDB global tables are a fully managed, multi-Region, and multi-active database option that delivers fast and localized read and write performance for massively scaled global applications.
upvoted 1 times

Dantecito 3 months, 3 weeks ago

Selected Answer: A

A is the only option that allows you to write from any region. I understand the < 1 second latency, but DynamoDB global tables use a last-writer-wins reconciliation between concurrent updates.
B and C only supports read scaling, not low-latency global writes. It will be much much much much than 1 second to write on the database.
D I stopped reading at lambda so NO!.
upvoted 1 times

zdi561 4 months ago

Selected Answer: A

Dynamodb is consistent and read/write in all node, Aurora db , only the master is writable. For reservation , writing is criticle.
upvoted 1 times

khatingarun 4 months, 2 weeks ago

Selected Answer: A

A is the correct answer
upvoted 1 times

zdi561 4 months, 2 weeks ago

Selected Answer: B

dynamodb is not globally consistent. DynamoDB doesn't support strongly consistent reads across Regions.

https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/V2globaltables_HowItWorks.html

upvoted 2 times

hilker1983 5 months, 1 week ago

Selected Answer: B

correct answer bbbb

upvoted 1 times

EllenLiu 5 months, 3 weeks ago

Selected Answer: A

A: dynamoDB global table preview strong consistent

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/multi-region-strong-consistency-gt.html>

upvoted 1 times

Penjerla 5 months, 4 weeks ago

Selected Answer: A

Aurora replicas have to be in the same region as primary

<https://docs.aws.amazon.com/prescriptive-guidance/latest/aurora-replication-options/compare-solutions.html>

upvoted 2 times

FlyingHawk 5 months, 4 weeks ago

Selected Answer: B

Since DynamoDB Global Tables allow writes in all Regions, conflicts can occur if the same item is updated simultaneously in multiple Regions. DynamoDB uses a "last writer wins" conflict resolution strategy, which may lead to data integrity issues:

https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/V2globaltables_HowItWorks.html#V2globaltables_HowItWorks.conflict-resolution

If you require less than 1-second latency for write updates across Regions, DynamoDB may not consistently meet this requirement due to the asynchronous nature of its replication.

For this scenario, DynamoDB Global Tables do not fully meet the requirements due to their reliance on eventual consistency. If strong consistency is a must, go with Amazon Aurora Global Database.

upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: B

The question asks "Average latency must be less than 1 second on updates to the reservation database."

A is incorrect:

"Changes to a DynamoDB global tables are replicated asynchronously, with typical latency of between 0.5 - 2.5 seconds between AWS Regions in the same geographic area."

B is the answer:

"All Aurora Replicas return the same data for query results with minimal replica lag. This lag is usually much less than 100 milliseconds after the primary instance has written an update."

upvoted 2 times

MatAlves 8 months, 3 weeks ago

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.Replication.html>

<https://community.aws/content/2drxEND7MtOb2bWs2J0NICGewP/ddb-globaltables-lag?lang=en>

upvoted 3 times

SVDK 1 year, 4 months ago

Selected Answer: A

How can you update your database in the different regions with read replicas? You need to be able to read and write to the database from the different regions.

upvoted 3 times

Milivoje 1 year, 5 months ago

Selected Answer: A

In my Opinion it is A. The reason is that Aurora Read Replicas support up to 5 Read replicas in different regions . We don't have

that limitation with Dynamo DB Global tables, hence I vote for A.

upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: B

Purely from the wording, seems B.

DynamoDB "usually within one second"

Aurora "usually less than one second"

Question asks for "less than one second" thus Aurora

upvoted 3 times

Salilgen 5 months, 2 weeks ago

Answer is A.

Option B is about cross-region Aurora replicas not about Aurora Global Database.

"Cross-region Aurora replicas have a lag that depends on transaction volume. Typically, a few seconds for most systems."

<https://docs.aws.amazon.com/prescriptive-guidance/latest/aurora-replication-options/compare-solutions.html>

"In a global table, a newly written item is typically propagated to all replica tables within a second."

<https://aws.amazon.com/dynamodb/global-tables/>

upvoted 1 times

pentium75 1 year, 5 months ago

We need "a single primary reservation database that is globally consistent" -> A is out (DynamoDB is eventually consistent with "last writer wins" and "usually" updates "within [not: less than] one second"). D is out because it mentions multiple databases (and RDS Event Streams to not guarantee the order of events).

C is out because RDS has higher replication delay, only Aurora can guarantee "less than one second". So we'd have "a single primary reservation database that is globally consistent" in one region, and we'd have read replicas with "less than 1 second on updates" latency in other regions.

upvoted 5 times

numark 1 year, 6 months ago

"a web application for travel ticketing". This would be a transaction, so DynamoDB is not the answer.

upvoted 1 times

pentium75 1 year, 5 months ago

So you can't write to DynamoDB tables at all because tables writes are transactions?

upvoted 3 times

awsgeek75 1 year, 5 months ago

There are no assumptions about the application here. The choices are related to the database that has one primary source of truth but multi-region presence. No requirement for transaction is given or implied.

upvoted 2 times

Goutham4981 1 year, 6 months ago

Selected Answer: A

Dynamo DB global table acts as a single table. It does not consist of primary and standby databases. It is one single global table which is synchronously updated. Users can write to any of the regional endpoints and the write will be automatically updated across regions. To have a single primary database that is consistent does not align with dynamo db global tables.

Option B is even more dumb compared to A since read replicas does not provide failover capability or fast updates from the primary database.

The answer almost close to the requirement is Option A even though it is a misfit

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #508

Topic 1

A company has migrated multiple Microsoft Windows Server workloads to Amazon EC2 instances that run in the us-west-1 Region. The company manually backs up the workloads to create an image as needed.

In the event of a natural disaster in the us-west-1 Region, the company wants to recover workloads quickly in the us-west-2 Region. The company wants no more than 24 hours of data loss on the EC2 instances. The company also wants to automate any backups of the EC2 instances.

Which solutions will meet these requirements with the LEAST administrative effort? (Choose two.)

A. Create an Amazon EC2-backed Amazon Machine Image (AMI) lifecycle policy to create a backup based on tags. Schedule the backup to run twice daily. Copy the image on demand.

B. Create an Amazon EC2-backed Amazon Machine Image (AMI) lifecycle policy to create a backup based on tags. Schedule the backup to run twice daily. Configure the copy to the us-west-2 Region. **Most Voted**

C. Create backup vaults in us-west-1 and in us-west-2 by using AWS Backup. Create a backup plan for the EC2 instances based on tag values. Create an AWS Lambda function to run as a scheduled job to copy the backup data to us-west-2.

D. Create a backup vault by using AWS Backup. Use AWS Backup to create a backup plan for the EC2 instances based on tag values. Define the destination for the copy as us-west-2. Specify the backup schedule to run twice daily. **Most Voted**

E. Create a backup vault by using AWS Backup. Use AWS Backup to create a backup plan for the EC2 instances based on tag values. Specify the backup schedule to run twice daily. Copy on demand to us-west-2.

Correct Answer: BD

Community vote distribution

BD (100%)

Comments

cloudenthusiast **Highly Voted** 1 year, 6 months ago

Selected Answer: BD

Option B suggests using an EC2-backed Amazon Machine Image (AMI) lifecycle policy to automate the backup process. By configuring the policy to run twice daily and specifying the copy to the us-west-2 Region, the company can ensure regular backups are created and copied to the alternate region.

Option D proposes using AWS Backup, which provides a centralized backup management solution. By creating a backup vault and backup plan based on tag values, the company can automate the backup process for the EC2 instances. The backup schedule can be set to run twice daily, and the destination for the copy can be defined as the us-west-2 Region.

upvoted 10 times

cloudenthusiast 1 year, 6 months ago

Both options automate the backup process and include copying the backups to the us-west-2 Region, ensuring data resilience in the event of a disaster. These solutions minimize administrative effort by leveraging automated backup and copy mechanisms provided by AWS services.

upvoted 7 times

awsgeek75 Highly Voted 10 months, 3 weeks ago

Selected Answer: BD

LEAST admin overhead:

A: On demand so wrong

C: Lambda is overhead

E: On-demand is wrong

BD is the only choice. Although D seems to cover for B also, happy to be corrected.

upvoted 7 times

pmlabs Most Recent 1 year, 2 months ago

B D seems to meet the requiremnts fully

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: BD

B and D are the options that meet the requirements with the least administrative effort.

B uses EC2 image lifecycle policies to automatically create AMLs of the instances twice daily and copy them to the us-west-2 region. This automates regional backups.

D leverages AWS Backup to define a backup plan that runs twice daily and copies backups to us-west-2. AWS Backup automates EC2 instance backups.

Together, these options provide automated, regional EC2 backup capabilities with minimal administrative overhead.

upvoted 2 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: BD

options B and D will provide least administrative effort.

upvoted 2 times

antropaws 1 year, 6 months ago

Selected Answer: BD

I also vote B and D.

upvoted 2 times

nosense 1 year, 6 months ago

Selected Answer: BD

solutions are both automated and require no manual intervention to create or copy backups

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #509


Topic 1

A company operates a two-tier application for image processing. The application uses two Availability Zones, each with one public subnet and one private subnet. An Application Load Balancer (ALB) for the web tier uses the public subnets. Amazon EC2 instances for the application tier use the private subnets.

Users report that the application is running more slowly than expected. A security audit of the web server log files shows that the application is receiving millions of illegitimate requests from a small number of IP addresses. A solutions architect needs to resolve the immediate performance problem while the company investigates a more permanent solution.

What should the solutions architect recommend to meet this requirement?

- A. Modify the inbound security group for the web tier. Add a deny rule for the IP addresses that are consuming resources.
- B. Modify the network ACL for the web tier subnets. Add an inbound deny rule for the IP addresses that are consuming resources. **Most Voted**
- C. Modify the inbound security group for the application tier. Add a deny rule for the IP addresses that are consuming resources.
- D. Modify the network ACL for the application tier subnets. Add an inbound deny rule for the IP addresses that are consuming resources.

Correct Answer: B*Community vote distribution***Comments****lucdt4** **Highly Voted** 1 year, 6 months ago**Selected Answer: B**

A wrong because security group can't deny (only allow)
upvoted 26 times

cloudenthusiast **Highly Voted** 1 year, 6 months ago**Selected Answer: B**

In this scenario, the security audit reveals that the application is receiving millions of illegitimate requests from a small number

of IP addresses. To address this issue, it is recommended to modify the network ACL (Access Control List) for the web tier subnets.

By adding an inbound deny rule specifically targeting the IP addresses that are consuming resources, the network ACL can block the illegitimate traffic at the subnet level before it reaches the web servers. This will help alleviate the excessive load on the web tier and improve the application's performance.

upvoted 10 times

awsgeek75 Most Recent 10 months, 3 weeks ago

Selected Answer: B

A: Wrong as SG cannot deny. By default everything is deny in SG and you allow stuff

CD: App tier is not under attack so these are irrelevant options

B: Correct as NACL is exactly for this access control list to define rules for CIDR or IP addresses

upvoted 3 times

Besisco 4 months ago

agree, even if an app tier was under attack, the block rule should be applied at a web tier ACL

upvoted 1 times

TariqKipkemei 1 year, 1 month ago

Selected Answer: B

Modify the network ACL for the web tier subnets. Add an inbound deny rule for the IP addresses that are consuming resources.

upvoted 3 times

potomac 1 year, 1 month ago

Selected Answer: B

A is wrong

Security groups act at the network interface level, not the subnet level, and they support Allow rules only.

upvoted 3 times

Devsin2000 1 year, 2 months ago

Selected Answer: A

The security Group can be applied to an ALB at web tier.

upvoted 1 times

OSHOAIB 11 months ago

Security group rules are always permissive; you can't create rules that deny access.

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/security-group-rules.html>

upvoted 3 times

Goutham4981 1 year ago

Security group can't deny.

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Since the bad requests are targeting the web tier, adding ACL deny rules for those IP addresses on the web subnets will block the traffic before it reaches the instances.

Security group changes (Options A and C) would not be effective since the requests are not even reaching those resources.

Modifying the application tier ACL (Option D) would not stop the bad traffic from hitting the web tier.

upvoted 3 times

fakrap 1 year, 6 months ago

Selected Answer: B

A is wrong because you cannot put any deny in security group

upvoted 3 times

Rob1L 1 year, 6 months ago

Selected Answer: B

You cannot Deny on SG, so it's B
upvoted 6 times

nosense 1 year, 6 months ago

Selected Answer: A

Option B is not as effective as option A
upvoted 5 times

y0 1 year, 6 months ago

Security group only have allow rules
upvoted 3 times

nosense 1 year, 6 months ago

yeah, my mistake. B should be
upvoted 2 times

cloudenthusiast 1 year, 6 months ago

A and C out due to the fact that SG does not have deny on allow rules.
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #510

Topic 1

A global marketing company has applications that run in the ap-southeast-2 Region and the eu-west-1 Region. Applications that run in a VPC in eu-west-1 need to communicate securely with databases that run in a VPC in ap-southeast-2.

Which network design will meet these requirements?

- A. Create a VPC peering connection between the eu-west-1 VPC and the ap-southeast-2 VPC. Create an inbound rule in the eu-west-1 application security group that allows traffic from the database server IP addresses in the ap-southeast-2 security group.
- B. Configure a VPC peering connection between the ap-southeast-2 VPC and the eu-west-1 VPC. Update the subnet route tables. Create an inbound rule in the ap-southeast-2 database security group that references the security group ID of the application servers in eu-west-1.
- C. Configure a VPC peering connection between the ap-southeast-2 VPC and the eu-west-1 VPC. Update the subnet route tables. Create an inbound rule in the ap-southeast-2 database security group that allows traffic from the eu-west-1 application server IP addresses. **Most Voted**
- D. Create a transit gateway with a peering attachment between the eu-west-1 VPC and the ap-southeast-2 VPC. After the transit gateways are properly peered and routing is configured, create an inbound rule in the database security group that references the security group ID of the application servers in eu-west-1.

Correct Answer: C*Community vote distribution***Comments****VellaDevil** **Highly Voted** 1 year, 11 months ago**Selected Answer: C**

Answer: C --> "You cannot reference the security group of a peer VPC that's in a different Region. Instead, use the CIDR block of the peer VPC."

<https://docs.aws.amazon.com/vpc/latest/peering/vpc-peering-security-groups.html>

upvoted 49 times

MatAlves 8 months, 3 weeks ago

Wow, big thanks!
upvoted 2 times

hsinchang 1 year, 10 months ago

Thanks for this clarification!
upvoted 4 times

Axeashes Highly Voted 1 year, 11 months ago

Selected Answer: C

"You cannot reference the security group of a peer VPC that's in a different Region. Instead, use the CIDR block of the peer VPC."

<https://docs.aws.amazon.com/vpc/latest/peering/vpc-peering-security-groups.html>

upvoted 12 times

potomac Most Recent 1 year, 7 months ago

Selected Answer: C

After establishing the VPC peering connection, the subnet route tables need to be updated in both VPCs to route traffic to the other VPC's CIDR blocks through the peering connection.

upvoted 3 times

Bennyboy789 1 year, 9 months ago

Selected Answer: C

VPC Peering Connection: This allows communication between instances in different VPCs as if they are on the same network. It's a straightforward approach to connect the two VPCs.

Subnet Route Tables: After establishing the VPC peering connection, the subnet route tables need to be updated in both VPCs to route traffic to the other VPC's CIDR blocks through the peering connection.

Inbound Rule in Database Security Group: By creating an inbound rule in the ap-southeast-2 database security group that allows traffic from the eu-west-1 application server IP addresses, you ensure that only the specified application servers from the eu-west-1 VPC can access the database servers in the ap-southeast-2 VPC.

upvoted 4 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

B) Configure VPC peering between ap-southeast-2 and eu-west-1 VPCs. Update routes. Allow traffic in ap-southeast-2 database SG from eu-west-1 application server SG.

This option establishes the correct network connectivity for the applications in eu-west-1 to reach the databases in ap-southeast-2:

VPC peering connects the two VPCs across regions - [https://docs.aws.amazon.com/vpc/latest/peering/what-is-vpc-peering.html#:~:text=You%20can%20create%20a%20VPC,%2DRegion%20VPC%20peering%20connection\).](https://docs.aws.amazon.com/vpc/latest/peering/what-is-vpc-peering.html#:~:text=You%20can%20create%20a%20VPC,%2DRegion%20VPC%20peering%20connection).)

Updating route tables enables routing between the VPCs

Security group rule allowing traffic from eu-west-1 application server SG to ap-southeast-2 database SG secures connectivity

upvoted 2 times

awsgeek75 1 year, 4 months ago

No, you cannot use a SG reference from another region so last part "Create an inbound rule in the ap-southeast-2 database security group that references the security group ID of the application servers in eu-west-1" cannot be setup. This is why B is wrong.

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Options A, C, D have flaws:

Option A peer direction is wrong

Option C opens databases to application server IP addresses rather than SG

Option D uses transit gateway which is unnecessary for just two VPCs

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

Selected C but B can also work

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

I just tried from the the console, You can specify the name or ID of another security group in the same region. To specify a security group in another AWS account (EC2-Classic only), prefix it with the account ID and a forward slash, for example: 111122223333/OtherSecurityGroup.

You can Specify a single IP address, or an IP address range in CIDR notation in the same/other region.

In the exam both option B and C would be a pass. In the real world both option will work.

upvoted 3 times

TariqKipkemei 1 year, 7 months ago

Correction, You cannot reference the security group of a peer VPC that's in a different Region. Instead, use the CIDR block of the peer VPC.

The C is the only option here.

<https://docs.aws.amazon.com/vpc/latest/peering/vpc-peering-security-groups.html#:~:text=You%20cannot-,reference,-the%20security%20group>

upvoted 4 times

awsgeek75 1 year, 4 months ago

This is why B is wrong. You can never access cross region security group id

upvoted 1 times

Chris22usa 1 year, 11 months ago

I realize D is right as ChatGpt indicates. Because here is not a problem just one application in a VPC connection to another in different region. Actually there many applications in different VPCs in a region which need to connect any other application crossingly in other region. So two transit gateway need to installed in two regions for multiple to multiple VPCs connections.

upvoted 2 times

Iragmt 1 year, 11 months ago

However, there was also a part of "create an inbound rule in the database security group that references the security group ID of the application servers in eu-west-1"

therefore, still C because we cannot reference SG ID of diff VPC, we should use the CIDR block

upvoted 1 times

Chris22usa 1 year, 11 months ago

post it on ChaptGpt and it give me answer D. what heck with this?

upvoted 2 times

haoAWS 1 year, 11 months ago

Selected Answer: C

B is wrong because It is in a different region, so reference to the security group ID will not work. A is wrong because you need to update the route table. The answer should be C.

upvoted 1 times

mattcl 1 year, 11 months ago

is B. what happens if application server IP addresses changes (Option C). You must change manually the IP in the security group again.

upvoted 1 times

antropaws 1 year, 11 months ago

Selected Answer: C

I thought B, but I vote C after checking Axeashes response.

upvoted 1 times

MalioMato 2 years ago

henoneto 2 years ago

Selected Answer: C

I think the answer is C because the security groups are in different VPCs. When the question wants to allow traffic from app vpc to database vpc i think using peering connection you will be able to add the security groups rules using private ip addresses of app servers. I don't think the database VPC will identify the security group id of another VPC.

upvoted 2 times

REzirezi 2 years ago

D You cannot create a VPC peering connection between VPCs in different regions.

upvoted 3 times

[Removed] 2 years ago

You can peer any two VPCs in different Regions, as long as they have distinct, non-overlapping CIDR blocks
<https://docs.aws.amazon.com/devicefarm/latest/developerguide/amazon-vpc-cross-region.html>

upvoted 3 times

fakrap 2 years ago

You can peer any two VPCs in different Regions, as long as they have distinct, non-overlapping CIDR blocks. This ensures that all of the private IP addresses are unique, and it allows all of the resources in the VPCs to address each other without the need for any form of network address translation (NAT).

upvoted 2 times

nosense 2 years ago

Selected Answer: B

b for me. bcs correct inbound rule, and not overhead

upvoted 2 times

cloudenthusiast 2 years ago

Selected Answer: B

Option B suggests configuring a VPC peering connection between the ap-southeast-2 VPC and the eu-west-1 VPC. By establishing this peering connection, the VPCs can communicate with each other over their private IP addresses.

Additionally, updating the subnet route tables is necessary to ensure that the traffic destined for the remote VPC is correctly routed through the VPC peering connection.

To secure the communication, an inbound rule is created in the ap-southeast-2 database security group. This rule references the security group ID of the application servers in the eu-west-1 VPC, allowing traffic only from those instances. This approach ensures that only the authorized application servers can access the databases in the ap-southeast-2 VPC.

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #511

Topic 1

A company is developing software that uses a PostgreSQL database schema. The company needs to configure multiple development environments and databases for the company's developers. On average, each development environment is used for half of the 8-hour workday.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure each development environment with its own Amazon Aurora PostgreSQL database
- B. Configure each development environment with its own Amazon RDS for PostgreSQL Single-AZ DB instances
- C. Configure each development environment with its own Amazon Aurora On-Demand PostgreSQL-Compatible database
Most Voted
- D. Configure each development environment with its own Amazon S3 bucket by using Amazon S3 Object Select

Correct Answer: C*Community vote distribution***Comments****cloudenthusiast** **Highly Voted** 2 years ago**Selected Answer: C**

Option C suggests using Amazon Aurora On-Demand PostgreSQL-Compatible databases for each development environment. This option provides the benefits of Amazon Aurora, which is a high-performance and scalable database engine, while allowing you to pay for usage on an on-demand basis. Amazon Aurora On-Demand instances are typically more cost-effective for individual development environments compared to the provisioned capacity options.

upvoted 12 times

cloudenthusiast 2 years ago

Option B suggests using Amazon RDS for PostgreSQL Single-AZ DB instances for each development environment. While Amazon RDS is a reliable and cost-effective option, it may have slightly higher costs compared to Amazon Aurora On-Demand instances.

upvoted 6 times

lragmt 1 year, 11 months ago

I'm thinking that it should be B, since question does not mention any requirement only cost effective and this is just an development environment I guess we can leverage the use of RDS free tier also

upvoted 3 times

Stranko Highly Voted 1 year, 3 months ago

Selected Answer: C

Guys, when you use the pricing calculator the cost between option B and C is really close. I doubt anyone wants to test on your knowledge of exact pricings in your region. I think that "On Demand" being explicitly specified in option C and not being specified in option B is the main difference here the exam wants to test. In that case I'd assume that option B means a constantly running instance and not "On Demand" which would make the choice pretty obvious. Again, I don't think AWS exam will test you on knowing that a single AZ is cheaper by 0,005 cents than Aurora :D

upvoted 11 times

tch Most Recent 2 months, 4 weeks ago

Selected Answer: C

There is no upfront commitment with Aurora. You pay an hourly charge for each instance that you launch, and when you're finished with an Aurora DB instance, you can delete it. You do not need to overprovision storage as a safety margin, and you only pay for the storage you actually consume- from AWS official link

upvoted 1 times

zdi561 4 months, 2 weeks ago

Selected Answer: B

B RDS single az is the cheapest. Aurora needs at least two azs.

upvoted 1 times

tonybuivannghia 7 months, 3 weeks ago

Selected Answer: C

I choose C because Aurora on-Demand is Aurora Serverless:

The Aurora Serverless is cost effective. Scale out fine-grained increments to provide just the right number of database resources and pay only for capacity consumed.

upvoted 2 times

a7md0 11 months, 1 week ago

Selected Answer: B

Single-AZ DB instances cheaper

upvoted 1 times

trinh_le 1 year, 1 month ago

Selected Answer: B

Single AZ more cost effective

upvoted 2 times

chasingsummer 1 year, 4 months ago

Selected Answer: B

1 instance(s) x 0.245 USD hourly x (4 / 24 hours in a day) x 730 hours in a month = 29.8083 USD ---> Amazon RDS PostgreSQL instances cost (monthly)

1 instance(s) x 0.26 USD hourly x (4 / 24 hours in a day) x 730 hours in a month = 31.6333 USD ---> Amazon Aurora PostgreSQL-Compatible DB instances cost (monthly)

upvoted 2 times

upliftinghut 1 year, 4 months ago

Selected Answer: C

C is correct because B is cheaper but they don't mention to stop the DB when not in use

upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

On-Demand is cheaper that Aurora or RDS because of low weekly usage

upvoted 2 times

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: C

We have environments that are used on average 4 hours per workday = 20 hours per week. So with option C (Aurora on-demand aka serverless) we pay for 20 hours per week. With option B (RDS) we pay for 168 hours per week (the answer does not mention anything about automating shutdown etc.).

So even if Aurora Serverless is slightly more expensive than RDS, C is cheaper because we pay only 20 (not 168) hours per week.

upvoted 4 times

Salilgen 5 months, 2 weeks ago

Question doesn't talk about serverless. Moreover, you can select on-demand pricing model for RDS for PostgreSQL Single-AZ too.

IMO answer is B

upvoted 1 times

Mikado211 1 year, 5 months ago

Selected Answer: B

Aurora on demand is (a little) more expensive than Aurora
Aurora is more expensive than RDS single instance

So cost effectiveness == RDS.

(B)

upvoted 1 times

pentium75 1 year, 5 months ago

But if you use the database only 20 hours per week (5 x 4), wouldn't you pay way less with Aurora serverless than with RDS?

upvoted 3 times

Murtadhaceit 1 year, 5 months ago

Selected Answer: B

AWS Services Calculator is showing B cheaper by less than a dollar for the same settings for both. I used "db.r6g.large" for RDS (Single-AZ) and Aurora and put 4 hours/day.

upvoted 7 times

Stranko 1 year, 3 months ago

I used the calculator, single AZ is cheaper for the exact same usage duration, if you pick On-Demand option for it too. In Aurora case (option C) you have "On Demand" explicitly specified, so if it has to be specified then I suppose that B option is about a constantly running instance. If B had an "On Demand" added, I'd vote B too.

upvoted 2 times

JoseVincent68 1 year, 6 months ago

Selected Answer: B

Amazon RDS Single AZ is cheaper than Aurora Multi-AZ

upvoted 1 times

Wayne23Fang 1 year, 7 months ago

Selected Answer: B

Aurora instances will cost you ~20% more than RDS MySQL Given the running hours the same.
Also Aurora is HA.

upvoted 1 times

baba365 1 year, 8 months ago

... just trying to trick you. Aurora on demand is Aurora Serverless.

upvoted 6 times

Anmol_1010 1 year, 7 months ago

that is good piece of infoamtion
upvoted 1 times

deechean 1 year, 9 months ago

Selected Answer: C

Aurora allows you to pay for the hours used. 4 hour every day, you only need 1/6 cost of 24 hours per day. You can check the Aurora pricing calculator.
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #512

Topic 1

A company uses AWS Organizations with resources tagged by account. The company also uses AWS Backup to back up its AWS infrastructure resources. The company needs to back up all AWS resources.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use AWS Config to identify all untagged resources. Tag the identified resources programmatically. Use tags in the backup plan. **Most Voted**
- B. Use AWS Config to identify all resources that are not running. Add those resources to the backup vault.
- C. Require all AWS account owners to review their resources to identify the resources that need to be backed up.
- D. Use Amazon Inspector to identify all noncompliant resources.

Correct Answer: A

Community vote distribution

A (100%)

Comments

cloudenthusiast **Highly Voted** 2 years ago

Selected Answer: A

This solution allows you to leverage AWS Config to identify any untagged resources within your AWS Organizations accounts. Once identified, you can programmatically apply the necessary tags to indicate the backup requirements for each resource. By using tags in the backup plan configuration, you can ensure that only the tagged resources are included in the backup process, reducing operational overhead and ensuring all necessary resources are backed up.

upvoted 5 times

Gape4 **Most Recent** 11 months, 2 weeks ago

Selected Answer: A

I will go for A. C and D doesn't make sense. B- resources not running? No

upvoted 2 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: A

Use AWS config to deploy the tag rule and remediate resources that are not compliant.

upvoted 4 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

This option has the least operational overhead:

AWS Config continuously evaluates resource configurations and can identify untagged resources

Resources can be programmatically tagged via the AWS SDK based on Config data

Backup plans can use tag criteria to automatically back up newly tagged resources

No manual review or resource discovery needed

upvoted 3 times

Bill1000 2 years ago

Selected Answer: A

Vote A

upvoted 3 times

nosense 2 years ago

Selected Answer: A

a valid for me

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #513

Topic 1

A social media company wants to allow its users to upload images in an application that is hosted in the AWS Cloud. The company needs a solution that automatically resizes the images so that the images can be displayed on multiple device types. The application experiences unpredictable traffic patterns throughout the day. The company is seeking a highly available solution that maximizes scalability.

What should a solutions architect do to meet these requirements?

- A. Create a static website hosted in Amazon S3 that invokes AWS Lambda functions to resize the images and store the images in an Amazon S3 bucket. **Most Voted**
- B. Create a static website hosted in Amazon CloudFront that invokes AWS Step Functions to resize the images and store the images in an Amazon RDS database.
- C. Create a dynamic website hosted on a web server that runs on an Amazon EC2 instance. Configure a process that runs on the EC2 instance to resize the images and store the images in an Amazon S3 bucket.
- D. Create a dynamic website hosted on an automatically scaling Amazon Elastic Container Service (Amazon ECS) cluster that creates a resize job in Amazon Simple Queue Service (Amazon SQS). Set up an image-resizing program that runs on an Amazon EC2 instance to process the resize jobs.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**cloudenthusiast** **Highly Voted** 2 years ago**Selected Answer: A**

By using Amazon S3 and AWS Lambda together, you can create a serverless architecture that provides highly scalable and available image resizing capabilities. Here's how the solution would work:

Set up an Amazon S3 bucket to store the original images uploaded by users. Configure an event trigger on the S3 bucket to invoke an AWS Lambda function whenever a new image is uploaded. The Lambda function can be designed to retrieve the uploaded image, perform the necessary resizing operations based on device requirements, and store the resized images back in the S3 bucket or a different bucket designated for resized images. Configure the Amazon S3 bucket to make the resized images publicly accessible for serving to users.

invoted 18 times

upvoted 10 times

surajkrishnamurthy Most Recent 4 months ago

Selected Answer: A

Is Static Website Read only ?
Eventhough I have selected Answer as A
upvoted 1 times

LeonSauveterre 6 months ago

Selected Answer: A

OK. A is better, although I chose D originally.

I guess it shouldn't be D simply because it overkills for this scenario and introduces higher operational overhead (even though option D can handle more complex workloads than Lambda if needed).

upvoted 1 times

cnureddy 1 year ago

How can end user upload an image to S3 bucket with static hosting. I believe it should be dynamic website (Answer D)

upvoted 2 times

mr123dd 1 year, 5 months ago

image = static = S3 or cloudfront
but image is unstructured data so you dont store it in a relational database like RDS
and Step Function is not for processing
So A

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

This meets all the key requirements:

S3 static website provides high availability and auto scaling to handle unpredictable traffic
Lambda functions invoked from the S3 site can resize images on the fly
Storing images in S3 buckets provides durability, scalability and high throughput
Serverless approach with S3 and Lambda maximizes scalability and availability

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: A

Scalability = S3, Lamda
automatically resize images = Lambda

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #514

Topic 1


A company is running a microservices application on Amazon EC2 instances. The company wants to migrate the application to an Amazon Elastic Kubernetes Service (Amazon EKS) cluster for scalability. The company must configure the Amazon EKS control plane with endpoint private access set to true and endpoint public access set to false to maintain security compliance. The company must also put the data plane in private subnets. However, the company has received error notifications because the node cannot join the cluster.

Which solution will allow the node to join the cluster?

- A. Grant the required permission in AWS Identity and Access Management (IAM) to the AmazonEKSNodeRole IAM role.
- B. Create interface VPC endpoints to allow nodes to access the control plane. **Most Voted**
- C. Recreate nodes in the public subnet. Restrict security groups for EC2 nodes.
- D. Allow outbound traffic in the security group of the nodes.

Correct Answer: B

Community vote distribution



Comments

cloudenthusiast **Highly Voted** 2 years ago

Selected Answer: B

By creating interface VPC endpoints, you can enable the necessary communication between the Amazon EKS control plane and the nodes in private subnets. This solution ensures that the control plane maintains endpoint private access (set to true) and endpoint public access (set to false) for security compliance.

upvoted 21 times

y0 **Highly Voted** 2 years ago

Selected Answer: A

Check this : <https://docs.aws.amazon.com/eks/latest/userguide/create-node-role.html>

Also, EKS does not require VPC endpoints. This is not the right use case for EKS

upvoted 21 times

TwinSpark 1 year, 1 month ago

correct i was going for B, but A looks better.

<https://docs.aws.amazon.com/eks/latest/userguide/cluster-endpoint.html>

"When you enable endpoint private access for your cluster, Amazon EKS creates a Route 53 private hosted zone on your behalf and associates it with your cluster's VPC. This private hosted zone is managed by Amazon EKS, and it doesn't appear in your account's Route 53 resources. "

upvoted 1 times

h0ng97_spare_002 1 year, 2 months ago

<https://docs.aws.amazon.com/eks/latest/userguide/create-node-role.html#:~:text=Before,launched>

"Before you can launch nodes and register them into a cluster, you must create an IAM role for those nodes to use when they are launched."

upvoted 4 times

JA2018 6 months, 2 weeks ago

which implies AmazonEKSTaskRole IAM role had already been configured..... that leaves answer B as the only viable choice

upvoted 1 times

bora4motion Most Recent 1 month, 1 week ago

Selected Answer: B

to me it looks like a network problem, not a "rights" problem.

upvoted 1 times

Dantecito 3 months, 3 weeks ago

Selected Answer: B

B. "The company has received error notifications because the node cannot join the cluster" it's a network connectivity issue, not permissions.

upvoted 1 times

4729e6c 4 months, 1 week ago

Selected Answer: B

Since the company has configured the Amazon EKS control plane with private access enabled (true) and public access disabled (false), the EKS API server is only accessible through private network routes. The worker nodes, which are in private subnets, must be able to communicate with the EKS control plane for proper functionality.

When the nodes cannot join the cluster, the most common cause is that they are unable to reach the EKS API server. In this case, since the control plane is private-only, AWS recommends using interface VPC endpoints (AWS PrivateLink) to enable private communication between the nodes and the control plane.

upvoted 1 times

Salilgen 5 months, 2 weeks ago

Selected Answer: A

IMO answer is A.

In private cluster you need interface VPC endpoints to necessary connect to some AWS services

(<https://docs.aws.amazon.com/eks/latest/userguide/private-clusters.html>) but EKS node always connect to control plane by EKS owned ENI.

See this: it is very clear.

<https://keetmalin.medium.com/eks-cluster-network-architecture-for-worker-nodes-635e067c8c2a>

upvoted 2 times

LeonSauveterre 6 months ago

Selected Answer: B

The question is trying to tell us:

1. "private access = true" and "public access = false". So the control plane endpoint is private and only accessible from within the VPC.

2. Nodes (data plane) are in private subnets.

Option A is NECESSARY but lack of permissions would generally cause authorization errors, not connectivity errors (that lead to failure of joining the cluster). Only after you got this right, you would likely receive errors about joining. So apparently we have already configured auth correctly, meaning A is not the answer.

Option C exposes the nodes to the internet. Wrong.

Option D is important for nodes to communicate with AWS services (because of allowing outbound traffic), but it's not sufficient if the required VPC interface endpoints are not even there.

upvoted 1 times

a7md0 11 months, 1 week ago

Selected Answer: A

AmazonEKSNodeRole IAM role

<https://docs.aws.amazon.com/eks/latest/userguide/create-node-role.html>

upvoted 2 times

emakid 11 months, 2 weeks ago

Selected Answer: B

When Amazon EKS nodes cannot join the cluster, especially when the control plane is set to private access only, the issue typically revolves around networking and connectivity. When the EKS control plane is configured with private access only, the nodes must communicate with the control plane over private IP addresses. Creating VPC endpoints (specifically, `com.amazonaws.<region>.eks`) allows traffic between the EKS nodes and the control plane to be routed privately within the VPC, which resolves the connectivity issue.

upvoted 3 times

Gape4 11 months, 2 weeks ago

Selected Answer: B

I think is B.

upvoted 2 times

MandAsh 11 months, 3 weeks ago

Selected Answer: B

Error they have mentioned is at network level. They are not saying authorisation is failed rather node is unable to connect to cluster aka connectivity issue. So answer it must be B

upvoted 2 times

Rocconno 12 months ago

Selected Answer: B

<https://docs.aws.amazon.com/eks/latest/userguide/private-clusters.html>

"Any self-managed nodes must be deployed to subnets that have the VPC interface endpoints that you require. If you create a managed node group, the VPC interface endpoint security group must allow the CIDR for the subnets, or you must add the created node security group to the VPC interface endpoint security group."

upvoted 2 times

stalk98 1 year, 1 month ago

I Think is A

upvoted 1 times

trinh_le 1 year, 1 month ago

Selected Answer: B

B is good to go

upvoted 3 times

JackyCCK 1 year, 2 months ago

S3/DynamoDB - VPC endpoint, other service should use interface endpoint so B is incorrect

upvoted 2 times

bujuman 1 year, 2 months ago

Selected Answer: B

Because of these two assertions:

- Amazon EKS control plane with endpoint private access set to true and endpoint public access set to false to maintain security compliance.

/* The command must also set the state value to private endpoints

(The company must also put the data plane in private subnets.

The best answer is related to Networking, Private Subnets (EKS Ctr Plane is strictly private and Data Plane stick under private subnets) and not related to EKS autodeployment that sure need an IAM policy. So according to me, answer B is the best answer.

upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: A

Before can launch nodes and register nodes into a EKS cluster, must create an IAM role for those nodes to use when they are launched.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #515

Topic 1

A company is migrating an on-premises application to AWS. The company wants to use Amazon Redshift as a solution.

Which use cases are suitable for Amazon Redshift in this scenario? (Choose three.)

A. Supporting data APIs to access data with traditional, containerized, and event-driven applications

B. Supporting client-side and server-side encryption **Most Voted**

C. Building analytics workloads during specified hours and when the application is not active **Most Voted**


D. Caching data to reduce the pressure on the backend database

E. Scaling globally to support petabytes of data and tens of millions of requests per minute **Most Voted**

F. Creating a secondary replica of the cluster by using the AWS Management Console

Correct Answer: BCE

Community vote distribution

**Comments**

elmogy **Highly Voted** 2 years ago

Selected Answer: BCE

Amazon Redshift is a data warehouse solution, so it is suitable for:

- Supporting encryption (client-side and server-side)
- Handling analytics workloads, especially during off-peak hours when the application is less active
- Scaling to large amounts of data and high query volumes for analytics purposes

The following options are incorrect because:

- A) Data APIs are not typically used with Redshift. It is more for running SQL queries and analytics.
- D) Redshift is not typically used for caching data. It is for analytics and data warehouse purposes.
- F) Redshift clusters do not create replicas in the management console. They are standalone clusters. you could create DR cluster from snapshot and restore to another region (automated or manual) but I do not think this what is meant in this option.

upvoted 21 times

pentium75 1 year, 5 months ago

"Data APIs are not typically used with Redshift" -> "With the Data API, you can programmatically access data in your Amazon Redshift cluster from different AWS services such as AWS Lambda, Amazon SageMaker notebooks, AWS Cloud9, and also your

on-premises applications using the AWS SDK. This allows you to build cloud-native, containerized, serverless, web-based, and event-driven applications on the AWS Cloud."

upvoted 2 times

Rob1L Highly Voted 2 years ago

Selected Answer: BCE

B. Supporting client-side and server-side encryption: Amazon Redshift supports both client-side and server-side encryption for improved data security.

C. Building analytics workloads during specified hours and when the application is not active: Amazon Redshift is optimized for running complex analytic queries against very large datasets, making it a good choice for this use case.

E. Scaling globally to support petabytes of data and tens of millions of requests per minute: Amazon Redshift is designed to handle petabytes of data, and to deliver fast query and I/O performance for virtually any size dataset.

upvoted 6 times

Dz1990 Most Recent 2 weeks, 1 day ago

Selected Answer: ACE

Going with ACE here!

The key insight about option B: encryption is a feature, not a use case. The question asks what you do with Redshift, not what it has. It's like asking "Why buy a Tesla?" and answering "it has seatbelts" vs "to reduce carbon footprint." See the difference?

For option E - yeah, "millions of requests" sounds intense for a data warehouse, but Redshift legitimately scales to petabytes and handles massive analytical workloads. It's not about transactional speed; it's about analytical scale. With concurrency scaling, it can handle way more than people think.

ACE = actual business use cases: modern API access, off-peak analytics, and massive data scaling. BCE = mixing features with use cases, which is a classic AWS exam trap!

upvoted 1 times

Dantecito 3 months, 3 weeks ago

Selected Answer: ABC

Chatgpt says ABC.

D. Not for caching.

E. It is not optimized for high-volume like millions of requests per minute.

F. Does not natively support creating a secondary replica of the cluster.

upvoted 1 times

LeonSauveterre 6 months ago

Selected Answer: BCE

YOU MUST KNOW: Amazon Redshift is a cloud-based data warehouse service designed for analytics and reporting on large datasets, optimized for complex queries and analytical workloads rather than transactional or real-time applications.

After you memorized that, option A is incorrect. If you want to work on operational workloads, then Aurora, Amazon RDS, DynamoDB, and such, are better choices.

Option D: Caching is not a core use case for Redshift. You can choose DAX (Amazon DynamoDB Accelerator) or ElastiCache to mitigate database pressure.

Option E: I choose that. BUT, Redshift is not designed for *high-concurrency* transactional workloads like "10s of mils of reqs/min", although it CAN handle massive datasets. Well... It seems partially true to me compared to other definitely wrong options.

upvoted 2 times

rpmaws 9 months ago

Selected Answer: ACE

B is not correct, how it can do encryption at client side ?

upvoted 2 times

3bdf1cc 11 months, 1 week ago

Found this related to A -- but specific to Redshift Serverless - but should qualify as a Redshift use case

The Data API enables you to seamlessly access data from Redshift Serverless with all types of traditional, cloud-native, and containerized serverless web service-based applications and event-driven applications.

<https://www.amazonaws.cn/en/blog-selection/use-the-amazon-redshift-data-api-to-interact-with-amazon-redshift-serverless/>

upvoted 1 times

NSA_Poker 1 year ago

Selected Answer: BCE

The following are obviously incorrect:

(D) Redshift is not as suitable as ElastiCache for caching.

(F) A secondary replica of the cluster is not supported.

The debate is between BCE & ACE or simplified, between A & C.

(A) is incorrect bc there is a difference btw Amazon Redshift Data API & API Gateway. API Gateway supports containerized and serverless workloads, as well as web applications. Amazon Redshift Data API is a built in API to access Redshift data with web services-based applications, including AWS Lambda, Amazon SageMaker notebooks, and AWS Cloud9.

<https://aws.amazon.com/blogs/big-data/build-a-serverless-analytics-application-with-amazon-redshift-and-amazon-api-gateway/>

(B) is correct. You have the following options of protecting data at rest in Amazon Redshift. Use server-side encryption OR use client-side encryption

<https://docs.aws.amazon.com/redshift/latest/mgmt/security-encryption.html>

upvoted 2 times

JackyCCK 1 year, 2 months ago

Redshift is OLAP(online analytical processing) so D is wrong, "when the application is not active"

upvoted 2 times

JackyCCK 1 year, 2 months ago

*C is wrong

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: ACE

A, C, E are for data and Redshift is data warehouse.

B is too generic of a choice

D caching is not the main purpose of Redshift

F replication is not main use of Redshift

CE are easy

Between AB, I chose A because Redshift supports data API and client-side encryption is not Redshift specific

upvoted 3 times

1rob 1 year, 5 months ago

Selected Answer: ABD

A: source <https://aws.amazon.com/blogs/big-data/using-the-amazon-redshift-data-api-to-interact-with-amazon-redshift-clusters/>

B: source: <https://docs.aws.amazon.com/redshift/latest/mgmt/security-encryption.html>

C: not sure, but you can configure scheduled queries, but the remark " and when the application is not active " , that is not relevant.

D: source https://docs.aws.amazon.com/redshift/latest/dg/c_challenges_achieving_high_performance_queries.html

E: Scaling globally is not supported; redshift is only a regional service.

F: only read replica is supported. So not a secondary replica of the cluster.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: ABD

A: <https://aws.amazon.com/de/blogs/big-data/get-started-with-the-amazon-redshift-data-api/>

B: <https://docs.aws.amazon.com/redshift/latest/mgmt/security-encryption.html>

D: https://docs.aws.amazon.com/redshift/latest/dg/c_challenges_achieving_high_performance_queries.html#result-caching

Not C: Redshift is a Data Warehouse; you can use that for analytics, but it is not directly related to an "application"

Not E: "Petabytes of data" yes, but "tens of millions of requests per minute" is not a typical feature of Redshift

Nor F: Replicas are not a Redshift feature

upvoted 2 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: ACE

Technically both options A and B apply, this is from the links below:

A. You can access your Amazon Redshift database using the built-in Amazon Redshift Data API.

<https://docs.aws.amazon.com/redshift/latest/mgmt/data-api.html#:~:text=in%20Amazon%20Redshift-,Data%20API,-. %20Using%20this%20API>

B. You can encrypt data client-side and upload the encrypted data to Amazon Redshift. In this case, you manage the encryption process, the encryption keys, and related tools.

<https://docs.aws.amazon.com/redshift/latest/mgmt/security-encryption.html#:~:text=Use-,client%2Dside,- encryption%20%E2%80%93%20You%20can>

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: ABC

Amazon Redshift provides a Data API that you can use to painlessly access data from Amazon Redshift with all types of traditional, cloud-native, and containerized, serverless web services-based and event-driven applications.

Amazon Redshift supports up to 500 concurrent queries per cluster, which may be expanded by adding more nodes to the cluster.

upvoted 3 times

potomac 1 year, 7 months ago

change to ABD

To reduce query runtime and improve system performance, Amazon Redshift caches the results of certain types of queries in memory on the leader node. When a user submits a query, Amazon Redshift checks the results cache for a valid, cached copy of the query results. If a match is found in the result cache, Amazon Redshift uses the cached results and doesn't run the query. Result caching is transparent to the user.

upvoted 1 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: BCE

The key use cases for Amazon Redshift that fit this scenario are:

B) Redshift supports both client-side and server-side encryption to protect sensitive data.

C) Redshift is well suited for running batch analytics workloads during off-peak times without affecting OLTP systems.

E) Redshift can scale to massive datasets and concurrent users to support large analytics workloads.

upvoted 3 times

cd93 1 year, 9 months ago

Selected Answer: BCD

Why E lol? It's a data warehouse! it has no need to support millions of requests, it is not mentioned anywhere (<https://aws.amazon.com/redshift/features>)

In fact Redshift editor supports max 500 connections and workgroup support max 2000 connections at once, see it's quota page

Redshift has a cache layer, D is correct

upvoted 3 times

mrsoa 1 year, 10 months ago

Selected Answer: BCE

BCE, For B this is why

<https://docs.aws.amazon.com/redshift/latest/mgmt/security-encryption.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #516

Topic 1

A company provides an API interface to customers so the customers can retrieve their financial information. The company expects a larger number of requests during peak usage times of the year.


The company requires the API to respond consistently with low latency to ensure customer satisfaction. The company needs to provide a compute host for the API.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use an Application Load Balancer and Amazon Elastic Container Service (Amazon ECS).
- B. Use Amazon API Gateway and AWS Lambda functions with provisioned concurrency. **Most Voted**
- C. Use an Application Load Balancer and an Amazon Elastic Kubernetes Service (Amazon EKS) cluster.
- D. Use Amazon API Gateway and AWS Lambda functions with reserved concurrency.

Correct Answer: B

Community vote distribution



Comments

cloudenthusiast **Highly Voted** 2 years ago

Selected Answer: B

In the context of the given scenario, where the company wants low latency and consistent performance for their API during peak usage times, it would be more suitable to use provisioned concurrency. By allocating a specific number of concurrent executions, the company can ensure that there are enough function instances available to handle the expected load and minimize the impact of cold starts. This will result in lower latency and improved performance for the API.

upvoted 13 times

Bennyboy789 **Highly Voted** 1 year, 9 months ago

Selected Answer: B

Provisioned - minimizing cold starts and providing low latency.

upvoted 7 times

LeonSauveterre **Most Recent** 6 months ago

Selected Answer: B

Concurrency of lambda function, provisioned vs reserved:

Reserved Concurrency:

1. Desired instances can start without interference from other functions
2. Guarantees available execution capacity for critical functions

Provisioned Concurrency:

1. Also guarantees consistent performance
2. [!!] Allows you to pre-warm a specific number of function instances to reduce cold start latency
upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

<https://docs.aws.amazon.com/lambda/latest/dg/lambda-concurrency.html#reserved-and-provisioned>

Consistency decreases if you exceed your provisioned instance. Lets say you have 1000 (default) provisioned instances and the load is 1500. The new 500 will have to wait until the first 1000 concurrent calls finish. This is solved by increasing the provisioned concurrency to 1500.

upvoted 2 times

1rob 1 year, 6 months ago

Selected Answer: A

So I have my doubts here. The question also states ;"The company needs to provide a compute host for the API." Imho this implies to have some sort of physical host which has to be provided by the customer. Translating this further to aws this would mean an EC2 instance. And then when I would go for ECS in stead of EKS.

Please share your opinion.

upvoted 7 times

pdragon1981 1 year, 5 months ago

Exactly, initially I was thinking on B but if company must provide a host I would say that only option A is feasible

upvoted 3 times

pdragon1981 1 year, 5 months ago

Sorry I understand bad the text, coorect answer is B, as for my understanding now the host is the device that the costumer needs to connect with API Gateway, bellow explains well the logic

<https://aws.amazon.com/api-gateway/>

upvoted 4 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

This option provides the least operational overhead:

API Gateway handles the API requests and integration with Lambda

Lambda automatically scales compute without managing servers

Provisioned concurrency ensures consistent low latency by keeping functions initialized

No need to manage containers or orchestration platforms as with ECS/EKS

upvoted 3 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: B

The company requires the API to respond consistently with low latency to ensure customer satisfaction especially during high peak periods, there is no mention of cost efficient. Hence provisioned concurrency is the best option.

Provisioned concurrency is the number of pre-initialized execution environments you want to allocate to your function. These execution environments are prepared to respond immediately to incoming function requests. Configuring provisioned concurrency incurs charges to your AWS account.

<https://docs.aws.amazon.com/lambda/latest/dg/provisioned-concurrency.html#:~:text=for%20a%20function.-,Provisioned%20concurrency,-%E2%80%93%20Provisioned%20concurrency%20is>

upvoted 2 times

MirKhobaeb 2 years ago

Selected Answer: B

Selected Answer: B

AWS Lambda provides a highly scalable and distributed infrastructure that automatically manages the underlying compute resources. It automatically scales your API based on the incoming request load, allowing it to respond consistently with low latency, even during peak times. AWS Lambda takes care of infrastructure provisioning, scaling, and resource management, allowing you to focus on writing the code for your API logic.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #517

Topic 1

A company wants to send all AWS Systems Manager Session Manager logs to an Amazon S3 bucket for archival purposes.

Which solution will meet this requirement with the MOST operational efficiency?

- A. Enable S3 logging in the Systems Manager console. Choose an S3 bucket to send the session data to. **Most Voted**
- B. Install the Amazon CloudWatch agent. Push all logs to a CloudWatch log group. Export the logs to an S3 bucket from the group for archival purposes.
- C. Create a Systems Manager document to upload all server logs to a central S3 bucket. Use Amazon EventBridge to run the Systems Manager document against all servers that are in the account daily.
- D. Install an Amazon CloudWatch agent. Push all logs to a CloudWatch log group. Create a CloudWatch logs subscription that pushes any incoming log events to an Amazon Kinesis Data Firehose delivery stream. Set Amazon S3 as the destination.

Correct Answer: A*Community vote distribution***Comments****master9** **Highly Voted** 1 year, 5 months ago**Selected Answer: A**

send logs to Amazon S3 from AWS Systems Manager Session Manager. Here are the steps to do so:

Enable S3 Logging: Open the AWS Systems Manager console. In the navigation pane, choose Session Manager. Choose the Preferences tab, and then choose Edit. Select the check box next to Enable under S3 logging.

Create an S3 Bucket: To store the Session Manager logs, create an S3 bucket to hold the audit logs from the Session Manager interactive shell usage.

Configure IAM Role: AWS Systems Manager Agent (SSM Agent) uses the same AWS Identity and Access Management (IAM) role to activate itself and upload logs to Amazon S3. You can use either an IAM instance profile that's attached to an Amazon Elastic Compute Cloud (Amazon EC2) instance or the IAM role that's configured for the Default Host Management Configuration.

upvoted 8 times

cloudenthusiast **Highly Voted** 2 years ago

Selected Answer: A

option A does not involve CloudWatch, while option D does. Therefore, in terms of operational overhead, option A would generally have less complexity and operational overhead compared to option D.

Option A simply enables S3 logging in the Systems Manager console, allowing you to directly send session logs to an S3 bucket. This approach is straightforward and requires minimal configuration.

On the other hand, option D involves installing and configuring the Amazon CloudWatch agent, creating a CloudWatch log group, setting up a CloudWatch Logs subscription, and configuring an Amazon Kinesis Data Firehose delivery stream to store logs in an S3 bucket. This requires additional setup and management compared to option A.

So, if minimizing operational overhead is a priority, option A would be a simpler and more straightforward choice.

upvoted 5 times

pujithacg8 Most Recent 10 months, 1 week ago

A, You can choose to store session log data in a specified Amazon Simple Storage Service (Amazon S3) bucket for debugging and troubleshooting purposes.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html#session-manager-logging-s3>

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

Most efficient is A because it is a direct option in SM logging.

B can work but is more operational overhead as you end up using CloudWatch (not sure how but making assumption based on language of option)

C is definitely too much work

D Way too many moving parts

upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: A

You can choose to store session log data in a specified Amazon Simple Storage Service (Amazon S3) bucket for debugging and troubleshooting purposes.

upvoted 2 times

deechean 1 year, 9 months ago

Selected Answer: A

You can config the log archived to S3 in the Session Manager - > preference tab. Another option is CloudWatch log.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html#session-manager-logging-s3>

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

*Simplicity - Enabling S3 logging requires just a simple configuration in the Systems Manager console to specify the destination S3 bucket. No other services need to be configured.

*Direct integration - Systems Manager has native support to send session logs to S3 through this feature. No need for intermediary services.

*Automated flow - Once S3 logging is enabled, the session logs automatically flow to the S3 bucket without manual intervention.

*Easy management - The S3 bucket can be managed independently for log storage and archival purposes without impacting Systems Manager.

*Cost-effectiveness - No charges for intermediate CloudWatch or Kinesis services. Just basic S3 storage costs.

*Minimal overhead - No ongoing management of complex pipeline of services. Direct logs to S3 minimizes overhead.

upvoted 3 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: A

With the MOST operational efficiency then option A is best.

Otherwise B is also an option with a little bit more ops than option A.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html>

upvoted 2 times

upvoted 2 times

Zox42 1 year, 11 months ago

Selected Answer: A

Answer A. <https://aws-labs.net/winlab5-manageinfra/sessmgrlog.html>

upvoted 2 times

Zuit 1 year, 11 months ago

Selected Answer: A

GPT argued for D.

B could be an option, by installing a logging package on alle managed systems/ECs etc.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/distributor-working-with-packages-deploy.html>

However, as it mentions the "Session manager logs" I would tend towards A.

upvoted 2 times

MrAWSAssociate 1 year, 11 months ago

Selected Answer: A

It should be "A".

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html>

upvoted 2 times

secdgs 1 year, 11 months ago

Selected Answer: A

It have menu to Enable S3 Logging.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html#session-manager-logging-s3>

upvoted 2 times

Markie999 2 years ago

Selected Answer: B

BBBBBBBBBB

upvoted 1 times

pentium75 1 year, 5 months ago

"Install the CloudWatch agent" where?

upvoted 2 times

Bill1000 2 years ago

Selected Answer: B

The option 'A' says "Enable S3 logging in the Systems Manager console." This means that you will enable the logs !! FOR !! S3 events and its is not what the question asks. My vote is for Option B, based on this article:

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/logging-with-S3.html>

upvoted 1 times

baba365 1 year, 11 months ago

To log session data using Amazon S3 (console)

Open the AWS Systems Manager console at <https://console.aws.amazon.com/systems-manager/>.

In the navigation pane, choose Session Manager.

Choose the Preferences tab, and then choose Edit.

Select the check box next to Enable under S3 logging.

upvoted 3 times

vrevkov 1 year, 11 months ago

But where do you want to install the Amazon CloudWatch agent in case of B?

upvoted 1 times

omoakin 2 years ago

DDDDDD

upvoted 1 times

Anmol_1010 2 years ago

Option D is definitely not right,
Its option B

upvoted 2 times

omoakin 2 years ago

Chat GPT says option A is incorrect cos it requires enabling S3 logging in the system manager console only logs information about the systems manager service not the session logs
Says correct answer is B

upvoted 1 times

[Removed] 2 years ago

Question may not be very clear. A should be the answer. Below link is the documetation:

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager-logging.html#session-manager-logging-s3>

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #518

Topic 1

An application uses an Amazon RDS MySQL DB instance. The RDS database is becoming low on disk space. A solutions architect wants to increase the disk space without downtime.

Which solution meets these requirements with the LEAST amount of effort?

- A. Enable storage autoscaling in RDS **Most Voted**
- B. Increase the RDS database instance size
- C. Change the RDS database instance storage type to Provisioned IOPS
- D. Back up the RDS database, increase the storage capacity, restore the database, and stop the previous instance

Correct Answer: A

Community vote distribution

A (100%)

Comments

cloudenthusiast **Highly Voted** 2 years ago

Selected Answer: A

Enabling storage autoscaling allows RDS to automatically adjust the storage capacity based on the application's needs. When the storage usage exceeds a predefined threshold, RDS will automatically increase the allocated storage without requiring manual intervention or causing downtime. This ensures that the RDS database has sufficient disk space to handle the increasing storage requirements.

upvoted 12 times

Gape4 **Most Recent** 11 months, 2 weeks ago

Selected Answer: A

Autoscaling.... without downtime...

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: A

Amazon RDS for MariaDB, Amazon RDS for MySQL, Amazon RDS for PostgreSQL, Amazon RDS for SQL Server and Amazon RDS for Oracle support RDS Storage Auto Scaling. RDS Storage Auto Scaling automatically scales storage capacity in response to growing database workloads, with zero downtime.

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

This question is so obvious

upvoted 4 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: A

RDS Storage Auto Scaling continuously monitors actual storage consumption, and scales capacity up automatically when actual utilization approaches provisioned storage capacity. Auto Scaling works with new and existing database instances. You can enable Auto Scaling with just a few clicks in the AWS Management Console. There is no additional cost for RDS Storage Auto Scaling. You pay only for the RDS resources needed to run your applications.

<https://aws.amazon.com/about-aws/whats-new/2019/06/rds-storage-auto-scaling/#:~:text=of%20the%20rest-,RDS%20Storage%20Auto%20Scaling,-continuously%20monitors%20actual>

upvoted 3 times

james2033 1 year, 10 months ago

Selected Answer: A

Quote "Amazon RDS now supports Storage Auto Scaling" and "... with zero downtime." (Jun 20th 2019) at

<https://aws.amazon.com/about-aws/whats-new/2019/06/rds-storage-auto-scaling/>

upvoted 3 times

james2033 1 year, 10 months ago

Hello moderator, please help me delete this discussion, I already add content before this comment.

upvoted 1 times

james2033 1 year, 10 months ago

Selected Answer: A

See "Amazon RDS now supports Storage Auto Scaling. Posted On: Jun 20, 2019. Starting today, Amazon RDS for MariaDB, Amazon RDS for MySQL, Amazon RDS for PostgreSQL, Amazon RDS for SQL Server and Amazon RDS for Oracle support RDS Storage Auto Scaling. RDS Storage Auto Scaling automatically scales storage capacity in response to growing database workloads, with zero downtime." at <https://aws.amazon.com/about-aws/whats-new/2019/06/rds-storage-auto-scaling/>

upvoted 3 times

haoAWS 1 year, 11 months ago

Selected Answer: A

A is the best answer.

B will not work for increasing disk space, it only improve the IO performance.

C will not work because it will cause downtime.

D is too complicated and need much operational effort.

upvoted 2 times

[Removed] 2 years ago

<https://aws.amazon.com/about-aws/whats-new/2019/06/rds-storage-auto-scaling/>

upvoted 2 times

Anmol_1010 2 years ago

The key word is No Down time. A would be bewt option

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #519

Topic 1

A consulting company provides professional services to customers worldwide. The company provides solutions and tools for customers to expedite gathering and analyzing data on AWS. The company needs to centrally manage and deploy a common set of solutions and tools for customers to use for self-service purposes.

Which solution will meet these requirements?

- A. Create AWS CloudFormation templates for the customers.
- B. Create AWS Service Catalog products for the customers. **Most Voted**
- C. Create AWS Systems Manager templates for the customers.
- D. Create AWS Config items for the customers.

Correct Answer: B

Community vote distribution

B (100%)

Comments

cloudenthusiast **Highly Voted** 2 years ago

Selected Answer: B

AWS Service Catalog allows you to create and manage catalogs of IT services that can be deployed within your organization. With Service Catalog, you can define a standardized set of products (solutions and tools in this case) that customers can self-service provision. By creating Service Catalog products, you can control and enforce the deployment of approved and validated solutions and tools.

upvoted 10 times

Oblako 1 year, 6 months ago

"within your organization" => not for customers

upvoted 1 times

JA2018 6 months, 2 weeks ago

for this use case, in order to access the self-service functions, customers must have some form of AWS identities (e.g. accounts, users, user groups or roles) already setup in the service provider's AWS organisation.

upvoted 1 times

Guru4Cloud Highly Voted 1 year, 9 months ago

Selected Answer: B

Some key advantages of using Service Catalog:

Centralized management - Products can be maintained in a single catalog for easy discovery and governance.

Self-service access - Customers can deploy the solutions on their own without manual intervention.

Standardization - Products provide pre-defined templates for consistent deployment.

Access control - Granular permissions can be applied to restrict product visibility and access.

Reporting - Service Catalog provides detailed analytics on product usage and deployments.

upvoted 5 times

hsinchang Most Recent 1 year, 10 months ago

Selected Answer: B

CloudFormation: a code as infrastructure service

Systems Manager: management solution for resources

Config: assess, audit and evaluate configurations

Other options does not fit this scenario.

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: B

AWS Service Catalog lets you centrally manage your cloud resources to achieve governance at scale of your infrastructure as code (IaC) templates, written in CloudFormation or Terraform. With AWS Service Catalog, you can meet your compliance requirements while making sure your customers can quickly deploy the cloud resources they need.

<https://aws.amazon.com/servicecatalog/#:~:text=How%20it%20works-,AWS%20Service%20Catalog,-lets%20you%20centrally>

upvoted 2 times

Yadav_Sanjay 2 years ago

Selected Answer: B

<https://docs.aws.amazon.com/servicecatalog/latest/adminguide/introduction.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #520

Topic 1


A company is designing a new web application that will run on Amazon EC2 Instances. The application will use Amazon DynamoDB for backend data storage. The application traffic will be unpredictable. The company expects that the application read and write throughput to the database will be moderate to high. The company needs to scale in response to application traffic.

Which DynamoDB table configuration will meet these requirements MOST cost-effectively?

- A. Configure DynamoDB with provisioned read and write by using the DynamoDB Standard table class. Set DynamoDB auto scaling to a maximum defined capacity.
- B. Configure DynamoDB in on-demand mode by using the DynamoDB Standard table class. **Most Voted**
- C. Configure DynamoDB with provisioned read and write by using the DynamoDB Standard Infrequent Access (DynamoDB Standard-IA) table class. Set DynamoDB auto scaling to a maximum defined capacity.
- D. Configure DynamoDB in on-demand mode by using the DynamoDB Standard Infrequent Access (DynamoDB Standard-IA) table class.

Correct Answer: B

Community vote distribution



Comments

Efren **Highly Voted** 2 years ago

B for me. Provisioned if we know how much traffic will come, but its unpredictable, so we have to go for on-demand
upvoted 12 times

VellaDevil 1 year, 11 months ago

Spot On
upvoted 3 times

tch **Most Recent** 3 months ago

Selected Answer: B

DynamoDB offers two table classes designed to help you optimize for cost.

The DynamoDB Standard table class is the default, and is recommended for the vast majority of workloads. The DynamoDB Standard-Infrequent Access (DynamoDB Standard-IA) table class is optimized for tables where storage is the dominant cost. For example, tables that store infrequently accessed data, such as application logs, old social media posts, e-commerce order history, and past gaming achievements, are good candidates for the Standard-IA table class.

upvoted 1 times

fdae619 5 months, 2 weeks ago

Selected Answer: B

Can someone answer why the standard class is a better choice to standard IA?

upvoted 1 times

tch 3 months ago

DynamoDB offers two table classes designed to help you optimize for cost.

The DynamoDB Standard table class is the default, and is recommended for the vast majority of workloads.

The DynamoDB Standard-Infrequent Access (DynamoDB Standard-IA) table class is optimized for tables where storage is the dominant cost. For example, tables that store infrequently accessed data, such as application logs, old social media posts, e-commerce order history, and past gaming achievements, are good candidates for the Standard-IA table class.

upvoted 1 times

emakid 11 months, 2 weeks ago

Selected Answer: B

Configure DynamoDB in on-demand mode by using the DynamoDB Standard table class.

This option allows DynamoDB to automatically adjust to varying traffic patterns, which is ideal for unpredictable workloads. The Standard table class is suitable for applications with moderate to high read and write throughput, and on-demand mode ensures that you are billed based on the actual usage, providing cost efficiency for variable traffic patterns.

upvoted 3 times

Gape4 11 months, 2 weeks ago

Selected Answer: B

Key word : On demand. So I think B.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

On demand

<https://docs.aws.amazon.com/wellarchitected/latest/serverless-applications-lens/capacity.html>

"With on-demand capacity mode, DynamoDB charges you for the data reads and writes your application performs on your tables. You do not need to specify how much read and write throughput you expect your application to perform because DynamoDB instantly accommodates your workloads as they ramp up or down."

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: B

Not A because of "unpredictable" traffic

Not C and D because we are expecting "moderate to high" traffic

upvoted 4 times

leonliu4 1 year, 6 months ago

Selected Answer: B

Leaning towards B, it's hard to predict the capacity for A, and autoscaling doesn't respond fast

upvoted 3 times

peekingpicker 1 year, 6 months ago

Selected Answer: A

it's A.

remember that :

he company expects that the application read and write throughput to the database will be moderate to high

provisioned throughput is cheaper than ondemand capacity right ?

upvoted 2 times

pentium75 1 year, 5 months ago

but "unpredictable" which usually hints to on-demand

upvoted 2 times

dilaaziz 1 year, 6 months ago

Selected Answer: D

Data storage: <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithTables.tableclasses.html>

upvoted 1 times

Mr_Marcus 1 year ago

Actually, this link confirmed "B" for me...

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: B

On-demand mode is great for unpredictable traffic

upvoted 3 times

bsbs1234 1 year, 8 months ago

I choose B

I think the items stored in the table in this question has large size. So each read/write, a big chunk of data pass through. A capacity unit is used to describe data throughput. provision to the high capacity units will be a waste because unpredicted traffic pattern.

upvoted 2 times

Bennyboy789 1 year, 9 months ago

Selected Answer: B

Unpredictable= on demand

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

The key factors are:

With On-Demand mode, you only pay for what you use instead of over-provisioning capacity. This avoids idle capacity costs. DynamoDB Standard provides the fastest performance needed for moderate-high traffic apps vs Standard-IA which is for less frequent access.

Auto scaling with provisioned capacity can also work but requires more administrative effort to tune the scaling thresholds.

upvoted 2 times

msdnpro 1 year, 10 months ago

Selected Answer: B

Support for B from AWS:

On-demand mode is a good option if any of the following are true:

- You create new tables with unknown workloads.
- You have unpredictable application traffic.
- You prefer the ease of paying for only what you use.

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ReadWriteCapacityMode.html>

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: B

Technically both options A and B will work. But this statement 'traffic will be unpredictable' rules out option A, because 'provisioned mode' was made for scenarios where traffic is predictable.

So I will stick with B, because on-demand mode is made for unpredictable traffic and instantly accommodates workloads as

so I will stick with B, because on-demand mode is made for unpredictable traffic and instantly accommodates workloads as they ramp up or down.

upvoted 3 times

0628atv 1 year, 10 months ago

Selected Answer: A

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/AutoScaling.html>

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #521

Topic 1

A retail company has several businesses. The IT team for each business manages its own AWS account. Each team account is part of an organization in AWS Organizations. Each team monitors its product inventory levels in an Amazon DynamoDB table in the team's own AWS account.

The company is deploying a central inventory reporting application into a shared AWS account. The application must be able to read items from all the teams' DynamoDB tables.

Which authentication option will meet these requirements MOST securely?

- A. Integrate DynamoDB with AWS Secrets Manager in the inventory application account. Configure the application to use the correct secret from Secrets Manager to authenticate and read the DynamoDB table. Schedule secret rotation for every 30 days.
- B. In every business account, create an IAM user that has programmatic access. Configure the application to use the correct IAM user access key ID and secret access key to authenticate and read the DynamoDB table. Manually rotate IAM access keys every 30 days.
- C. In every business account, create an IAM role named BU_ROLE with a policy that gives the role access to the DynamoDB table and a trust policy to trust a specific role in the inventory application account. In the inventory account, create a role named APP_ROLE that allows access to the STS AssumeRole API operation. Configure the application to use APP_ROLE and assume the crossaccount role BU_ROLE to read the DynamoDB table. **Most Voted**
- D. Integrate DynamoDB with AWS Certificate Manager (ACM). Generate identity certificates to authenticate DynamoDB. Configure the application to use the correct certificate to authenticate and read the DynamoDB table.

Correct Answer: C

Community vote distribution

C (100%)

Comments

cloudenthusiast **Highly Voted** 2 years ago

Selected Answer: C

IAM Roles: IAM roles provide a secure way to grant permissions to entities within AWS. By creating an IAM role in each business account named BU_ROLE with the necessary permissions to access the DynamoDB table, the access can be controlled at the

account named BU_ROLE with the necessary permissions to access the DynamoDB table, the access can be controlled at the IAM role level.

Cross-Account Access: By configuring a trust policy in the BU_ROLE that trusts a specific role in the inventory application account (APP_ROLE), you establish a trusted relationship between the two accounts.

Least Privilege: By creating a specific IAM role (BU_ROLE) in each business account and granting it access only to the required DynamoDB table, you can ensure that each team's table is accessed with the least privilege principle.

Security Token Service (STS): The use of STS AssumeRole API operation in the inventory application account allows the application to assume the cross-account role (BU_ROLE) in each business account.

upvoted 31 times

TariqKipkemei 1 year, 10 months ago

Well broken down..thank you :)

upvoted 3 times

MandAsh Highly Voted 11 months, 3 weeks ago

C because they have taken effort to explain it in details.. lol

upvoted 6 times

Bennyboy789 Most Recent 1 year, 9 months ago

Selected Answer: C

Keyword: IAM ROLES

upvoted 4 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

C is the most secure option to meet the requirements.

Using cross-account IAM roles and role chaining allows the inventory application to securely access resources in other accounts. The roles provide temporary credentials and can be permissions controlled.

upvoted 3 times

hsinchang 1 year, 10 months ago

Selected Answer: C

Looks complex, but IAM role seems more probable, I go with C.

upvoted 4 times

mattcl 1 year, 11 months ago

Why not A?

upvoted 3 times

awsgeek75 1 year, 5 months ago

A is wrong because it is incomplete. Just integrating with secrets manager doesn't give any access to DynamoDB.

upvoted 2 times

antropaws 1 year, 11 months ago

Selected Answer: C

It's complex, but looks C.

upvoted 2 times

eehhssaan 2 years ago

i'll go with C .. coming from two minds

upvoted 3 times

nosense 2 years ago

a or c. C looks like a more secure

upvoted 2 times

omoakin 2 years ago

CCCCCCCCCCCC
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #522

Topic 1

A company runs container applications by using Amazon Elastic Kubernetes Service (Amazon EKS). The company's workload is not consistent throughout the day. The company wants Amazon EKS to scale in and out according to the workload.

Which combination of steps will meet these requirements with the LEAST operational overhead? (Choose two.)

A. Use an AWS Lambda function to resize the EKS cluster.

B. Use the Kubernetes Metrics Server to activate horizontal pod autoscaling. **Most Voted**

C. Use the Kubernetes Cluster Autoscaler to manage the number of nodes in the cluster. **Most Voted**

D. Use Amazon API Gateway and connect it to Amazon EKS.

E. Use AWS App Mesh to observe network activity.

Correct Answer: BC

Community vote distribution

BC (100%)

Comments

Guru4Cloud **Highly Voted** 9 months, 2 weeks ago

Selected Answer: BC

B and C are the correct options.

Using the Kubernetes Metrics Server (B) enables horizontal pod autoscaling to dynamically scale pods based on CPU/memory usage. This allows scaling at the application tier level.

The Kubernetes Cluster Autoscaler (C) automatically adjusts the number of nodes in the EKS cluster in response to pod resource requirements and events. This allows scaling at the infrastructure level.

upvoted 7 times

cloudenthusiast **Highly Voted** 1 year ago

Selected Answer: BC

By combining the Kubernetes Cluster Autoscaler (option C) to manage the number of nodes in the cluster and enabling horizontal pod autoscaling (option B) with the Kubernetes Metrics Server, you can achieve automatic scaling of your EKS cluster and container applications based on workload demand. This approach minimizes operational overhead as it leverages built-in Kubernetes functionality and automation mechanisms.

upvoted 5 times

wsdasdasdqwda **Most Recent** 7 months, 1 week ago

K8S Metrics Server and Autoscaler => B and C

upvoted 4 times

TariqKipkemei 10 months, 3 weeks ago

Selected Answer: BC

This is pretty straight forward.

Use the Kubernetes Metrics Server to activate horizontal pod autoscaling.

Use the Kubernetes Cluster Autoscaler to manage the number of nodes in the cluster.

upvoted 5 times

james2033 10 months, 3 weeks ago

Selected Answer: BC

Kubernetes Metrics Server <https://docs.aws.amazon.com/eks/latest/userguide/metrics-server.html>

AWS Autoscaler <https://docs.aws.amazon.com/eks/latest/userguide/autoscaling.html> and

<https://github.com/kubernetes/autoscaler/blob/master/cluster-autoscaler/cloudprovider/aws/README.md>

upvoted 4 times

nosense 1 year ago

Selected Answer: BC

b and c is right

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #523

Topic 1

A company runs a microservice-based serverless web application. The application must be able to retrieve data from multiple Amazon DynamoDB tables. A solutions architect needs to give the application the ability to retrieve the data with no impact on the baseline performance of the application.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. AWS AppSync pipeline resolvers
- B. Amazon CloudFront with Lambda@Edge functions
- C. Edge-optimized Amazon API Gateway with AWS Lambda functions

D. Amazon Athena Federated Query with a DynamoDB connector **Most Voted**

Correct Answer: D

Community vote distribution



Comments

elmogy **Highly Voted** 2 years ago

just passed yesterday 30-05-23, around 75% of the exam came from here, some with light changes.
upvoted 36 times

omoakin **Highly Voted** 2 years ago

Great work made it to the last question. Goodluck to you all
upvoted 20 times

cyber_bedouin 1 year, 7 months ago

Thanks. Do you think the questions after 500 are relevant, they seem to be above associate level (harder)
upvoted 10 times

Drew3000 1 year, 2 months ago

I am jealous that 10 months ago there were only 523 questions :(
upvoted 10 times

Awsbeginner87 1 year, 2 months ago

right...how to go through 840 questions :(
upvoted 3 times

Rhydian25 11 months, 2 weeks ago

Currently 904 :)
upvoted 5 times

4729e6c 4 months, 1 week ago

1019 now and hard to finish them all in a month :(
upvoted 3 times

bora4motion 1 month, 1 week ago

how did you do?
upvoted 1 times

Awsbeginner87 1 year, 2 months ago

moderators answer is wrong for many questions.which one should we consider:(
upvoted 1 times

MostofMichelle 2 years ago

good luck to you as well.
upvoted 6 times

CTao Most Recent 4 months ago

Selected Answer: A

AWS AppSync: It provides a managed GraphQL service that simplifies data access across multiple sources, including DynamoDB tables.
upvoted 2 times

4729e6c 4 months, 1 week ago

Selected Answer: A

Deepseek and Qwen are just published and they both say A.
upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

Selected Answer: A

AWS AppSync is a managed GraphQL service that allows applications to query multiple data sources, such as DynamoDB tables, in a single API call. It is highly efficient and does not require maintaining backend infrastructure.

Pipeline resolvers in AppSync enable the combination of multiple resolvers to fetch data from different sources, aggregate the data, and return it in a single response. This approach is operationally efficient, has minimal impact on baseline performance, and is designed for applications interacting with multiple DynamoDB tables.

<https://docs.aws.amazon.com/appsync/latest/devguide/tutorial-pipeline-resolvers.html>

upvoted 1 times

FlyingHawk 5 months, 4 weeks ago

Why the suggested answer is D? I read the comments, feel A is more suitable as the requirement is to give the application the ability to retrieve the data with no impact on the baseline performance of the application. D might have affect the baseline performance, A is more effective

upvoted 2 times

LeonSauveterre 6 months ago

Selected Answer: A

A - Correct. Retrieves data from multiple data sources.

B - Incorrect. Lambda@Edge extends serverless processing to CloudFront, but it is typically used for content manipulation, caching, or geographic distribution.

C - That works but requires a lot of things to do manually. Why C when you can do with A?

D - Amazon Athena Federated Query is designed for ad hoc querying of data across different sources BUT it has impact on performance, and not operationally efficient.

PS: "ad hoc" means "for this specific purpose". Such requests can lack efficiency for repetitive tasks because it's not optimized or automated, since they're naturally exploratory or tailored to specific immediate needs.

upvoted 2 times

kevindu 9 months, 3 weeks ago

Selected Answer: A

Is there anyone who has recently passed the exam who can tell me approximately how many of the original questions are in the actual exam?

upvoted 3 times

use4u 1 year, 1 month ago

let me know

for most of questions that system answer are correct or comment answer are correct?

upvoted 2 times

Sergiuss95 1 year, 1 month ago

Comment answer

upvoted 2 times

osmk 1 year, 3 months ago

Selected Answer: D

https://docs.amazonaws.cn/en_us/athena/latest/ug/connect-to-a-data-source.html

upvoted 6 times

FlyingHawk 4 months, 2 weeks ago

Athena queries are serverless and scalable but are not optimized for low-latency, high-concurrency scenarios typically required for a microservice-based operational web application. Additionally, query execution time and costs can become a concern if frequent access to DynamoDB tables is required.

upvoted 2 times

upliftinghut 1 year, 4 months ago

Selected Answer: D

key word is most operational effective => D requires no coding

upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

I'll go with D as ABC looks too much work or irrelevant. Although not sure how AFQ actually achieves the read without impacting performance.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: D

Not A - Pipe Resolvers require coding, would not consider that 'operationally efficient'

Not B - CloudFront caches web content at the edge, not DynamoDB query results for apps

Not C - Neither API Gateway or Lambda have anything to do with DynamoDB performance

D - Can do exactly that

upvoted 8 times

aws94 1 year, 5 months ago

Selected Answer: A

I am not an expert but I used Bing+Gemini+Chatgpt=AAA

upvoted 2 times

ekisako 1 year, 6 months ago

Selected Answer: A

multiple database tables = AppSync pipeline resolvers

upvoted 7 times

hungta 1 year, 6 months ago

Selected Answer: B

For an operationally efficient solution that minimizes impact on baseline performance in a microservice-based serverless web application retrieving data from multiple DynamoDB tables, Amazon CloudFront with Lambda@Edge functions (Option B) is often the most suitable choice

upvoted 1 times

pentium75 1 year, 5 months ago

CloudFront to retrieve data from DynamoDB tables?

upvoted 4 times

thanhnv142 1 year, 7 months ago

D is correct. There is construction of how to retrieve data from DynamoDB with Athena
<https://docs.aws.amazon.com/athena/latest/ug/connect-to-a-data-source.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #524

Topic 1

A company wants to analyze and troubleshoot Access Denied errors and Unauthorized errors that are related to IAM permissions. The company has AWS CloudTrail turned on.

Which solution will meet these requirements with the LEAST effort?

- A. Use AWS Glue and write custom scripts to query CloudTrail logs for the errors.
- B. Use AWS Batch and write custom scripts to query CloudTrail logs for the errors.
- C. Search CloudTrail logs with Amazon Athena queries to identify the errors. **Most Voted**
- D. Search CloudTrail logs with Amazon QuickSight. Create a dashboard to identify the errors.

Correct Answer: C*Community vote distribution***Comments****Guru4Cloud** **Highly Voted** 1 year, 3 months ago**Selected Answer: C**

Athena allows you to run SQL queries on data in Amazon S3, including CloudTrail logs. It is the easiest way to query the logs and identify specific errors without needing to write any custom code or scripts.

With Athena, you can write simple SQL queries to filter the CloudTrail logs for the "AccessDenied" and "UnauthorizedOperation" error codes. This will return the relevant log entries that you can then analyze.

upvoted 5 times

awsgeek75 **Highly Voted** 10 months, 3 weeks ago**Selected Answer: C**

<https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html>

When troubleshooting you will want to query specific things in the log and Athena provides query language for that.

Quick Sight is data analytics and visualisation tool. You can use it to aggregate data and maybe make a dashboard for number of errors by type etc but that doesn't help you troubleshoot anything.

C is correct

upvoted 5 times

pentium75 Most Recent 11 months, 1 week ago

Selected Answer: C

"Search CloudTrail logs with Amazon QuickSight", that doesn't work. QuickSight can visualize Athena query results, so "search CloudTrail logs with Amazon Athena, then create a dashboard with Amazon QuickSight" would make sense. But QuickSight without Athena won't work.

upvoted 4 times

[Removed] 12 months ago

Selected Answer: C

Athena is for searching

upvoted 3 times

bogobob 1 year ago

Selected Answer: D

The question asks specifically to "analyze and troubleshoot". While Athena is easy to get the data, you then just have a list of logs. Not very useful to troubleshoot...

upvoted 1 times

pentium75 11 months, 1 week ago

But without Athena, there is nothing you can visualize in QuickSight.

upvoted 3 times

awsgeek75 10 months, 3 weeks ago

How will pretty pictures in QuickSight help with troubleshooting?

upvoted 2 times

NickGordon 1 year, 1 month ago

Selected Answer: D

Quick Sight is an analytics tool. Sounds like a LEAST effort option

upvoted 3 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: C

C for me. Using Athena with CloudTrail logs is a powerful way to enhance your analysis of AWS service activity. For example, you can use queries to identify trends and further isolate activity by attributes, such as source IP address or user.

<https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html#:~:text=CloudTrail%20Lake%20documentation,-Using%20Athena,-with%20CloudTrail%20logs>

upvoted 2 times

james2033 1 year, 4 months ago

Selected Answer: C

IAM and CloudTrail <https://docs.aws.amazon.com/IAM/latest/UserGuide/cloudtrail-integration.html#stscldtrilexample-assumerole> .

Query CloudTrail logs by Athena [https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html#tips-for-querying-cloudtrail-logs](https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html#tips-for-querying-cloudtrail-logs#tips-for-querying-cloudtrail-logs)

upvoted 2 times

james2033 1 year, 4 months ago

Choose C, not D, because need "analyze and troubleshoot", not just see on dashboard (in D).

upvoted 2 times

live_reply_developers 1 year, 4 months ago

Selected Answer: C

Amazon Athena is an interactive query service provided by AWS that enables you to analyze data , is a little bit more suitable integrated with cloud trail that permit to verify WHO accessed the service.

upvoted 2 times

manuh 1 year, 5 months ago

Selected Answer: C

Dashboard isn't required. Also refer to this <https://repost.aws/knowledge-center/troubleshoot-iam-permission-errors>

upvoted 2 times

haoAWS 1 year, 5 months ago

Selected Answer: D

I am struggling for the C and D for a long time, and ask the chatGPT. The chatGPT says D is better, since Athena requires more expertise on SQL.

upvoted 1 times

antropaws 1 year, 5 months ago

Selected Answer: D

Both C and D are feasible. I vote for D:

Amazon QuickSight supports logging the following actions as events in CloudTrail log files:

- Whether the request was made with root or AWS Identity and Access Management user credentials
- Whether the request was made with temporary security credentials for an IAM role or federated user
- Whether the request was made by another AWS service

<https://docs.aws.amazon.com/quicksight/latest/user/logging-using-cloudtrail.html>

upvoted 1 times

PCWu 1 year, 5 months ago

Selected Answer: C

The Answer will be C:

Need to use Athena to query keywords and sort out the error logs.

D: No need to use Amazon QuickSight to create the dashboard.

upvoted 2 times

Axeashes 1 year, 5 months ago

Selected Answer: C

"Using Athena with CloudTrail logs is a powerful way to enhance your analysis of AWS service activity."

<https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html>

upvoted 2 times

oras2023 1 year, 6 months ago

Selected Answer: C

Analyse and TROUBLESHOOT, look like Athena

upvoted 1 times

oras2023 1 year, 5 months ago

<https://docs.aws.amazon.com/athena/latest/ug/cloudtrail-logs.html>

upvoted 2 times

alexandercamachop 1 year, 6 months ago

Selected Answer: D

It specifies analyze, not query logs.

Which is why option D is the best one as it provides dashboards to analyze the logs.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #525

Topic 1

A company wants to add its existing AWS usage cost to its operation cost dashboard. A solutions architect needs to recommend a solution that will give the company access to its usage cost programmatically. The company must be able to access cost data for the current year and forecast costs for the next 12 months.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Access usage cost-related data by using the AWS Cost Explorer API with pagination. **Most Voted**
- B. Access usage cost-related data by using downloadable AWS Cost Explorer report .csv files.
- C. Configure AWS Budgets actions to send usage cost data to the company through FTP.
- D. Create AWS Budgets reports for usage cost data. Send the data to the company through SMTP.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**BrijMohan08** **Highly Voted** 1 year, 3 months ago**Selected Answer: A**

Keyword

12 months, API Support

<https://docs.aws.amazon.com/cost-management/latest/userguide/ce-what-is.html>

upvoted 6 times

alexandercamachop **Highly Voted** 1 year, 6 months ago**Selected Answer: A**

Answer is: A

says dashboard = Cost Explorer, therefor C & D are eliminated.

also says programmatically, means non manual intervention therefor API.

upvoted 5 times

awsgeek75 **Most Recent** 11 months ago**Selected Answer: A**

Programmatically + LEAST overhead = API

programmatically + LEAST overhead = API
upvoted 4 times

TariqKipkemei 1 year ago

Selected Answer: A

access to its usage cost programmatically = AWS Cost Explorer API
upvoted 3 times

thanhnv142 1 year, 1 month ago

A: correct

1. programatically = API

2. In the next 12 months = cost explorer

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

Access usage cost-related data by using the AWS Cost Explorer API with pagination
upvoted 3 times

wendaz 1 year, 2 months ago

don't repeat the answer, it is useless... explain , okay? i have seen your replies many time just to copy the options.. it makes no sense...

upvoted 9 times

james2033 1 year, 4 months ago

Selected Answer: A

AWS Cost Explorer API with paginated request: <https://docs.aws.amazon.com/cost-management/latest/userguide/ce-api-best-practices.html#ce-api-best-practices-optimize-costs>

upvoted 3 times

MrAWSAssociate 1 year, 5 months ago

Selected Answer: A

From AWS Documentation*:

"You can view your costs and usage using the Cost Explorer user interface free of charge. You can also access your data programmatically using the Cost Explorer API. Each paginated API request incurs a charge of \$0.01. You can't disable Cost Explorer after you enable it."

* Source:

<https://docs.aws.amazon.com/cost-management/latest/userguide/ce-what-is.html>

<https://docs.aws.amazon.com/AWSJavaScriptSDK/v3/latest/clients/client-cost-explorer/interfaces/costexplorerpaginationconfiguration.html>

upvoted 5 times

oras2023 1 year, 6 months ago

Selected Answer: A

least operational overhead = API access
upvoted 4 times

oras2023 1 year, 6 months ago

least operational overhead = API access
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #526

Topic 1

A solutions architect is reviewing the resilience of an application. The solutions architect notices that a database administrator recently failed over the application's Amazon Aurora PostgreSQL database writer instance as part of a scaling exercise. The failover resulted in 3 minutes of downtime for the application.

Which solution will reduce the downtime for scaling exercises with the LEAST operational overhead?

- A. Create more Aurora PostgreSQL read replicas in the cluster to handle the load during failover.
- B. Set up a secondary Aurora PostgreSQL cluster in the same AWS Region. During failover, update the application to use the secondary cluster's writer endpoint.
- C. Create an Amazon ElastiCache for Memcached cluster to handle the load during failover.
- D. Set up an Amazon RDS proxy for the database. Update the application to use the proxy endpoint. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****alexandercamachop** **Highly Voted** 2 years ago**Selected Answer: D**

D is the correct answer.

It is talking about the write database. Not reader.

Amazon RDS proxy allows you to automatically route write request to the healthy writer, minimizing downtime.

upvoted 12 times

nilandd44gg 1 year, 10 months ago

One of the benefits of Amazon RDS Proxy is that it can improve application recovery time after database failovers. While RDS Proxy supports both MySQL as well as PostgreSQL engines, in this post, we will use a MySQL test workload to demonstrate how RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL and by up to 32% for Amazon RDS for MySQL.

<https://aws.amazon.com/blogs/database/improving-application-availability-with-amazon-rds-proxy/>

<https://aws.amazon.com/rds/proxy/faqs/>

upvoted 6 times

hora4motion **Most Recent** 1 month, 1 week ago

penetration 1 month, 1 week ago

Selected Answer: D

I'm going with D. It will reduce the failover time between DBs.

upvoted 1 times

ugaga 11 months, 3 weeks ago

AWS EXAM is also an AWS promotion...

RDS Proxy really does not reduce recovery time if you don't have alternative RDS instance. You will see if you understand what proxy is. But this is exam (= promotion) of AWS, D could be an answer, because some AWS document says so.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: D

"RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL "

<https://aws.amazon.com/de/blogs/database/improving-application-availability-with-amazon-rds-proxy/>

upvoted 3 times

ftaws 1 year, 5 months ago

Selected Answer: B

RDS Proxy is used for DB timeout not downtime.

How to reduce downtime with RDS Proxy?

There is no change downtime if we use RDS Proxy.

upvoted 3 times

pentium75 1 year, 5 months ago

"How to reduce downtime with RDS Proxy", by eliminating the need for the application to retrieve the new DNS record after the old one times out.

upvoted 4 times

MatAlves 8 months, 3 weeks ago

I would never have thought about DNS records in this question. Nice catch!

upvoted 2 times

pentium75 1 year, 5 months ago

"RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL "

<https://aws.amazon.com/de/blogs/database/improving-application-availability-with-amazon-rds-proxy/>

upvoted 3 times

Cyberkayu 1 year, 5 months ago

Selected Answer: B

they are using Aurora, RDS proxy dont work here

Answer B

upvoted 2 times

pentium75 1 year, 5 months ago

Wrong: <https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html>

upvoted 2 times

pentium75 1 year, 5 months ago

"RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL "

<https://aws.amazon.com/de/blogs/database/improving-application-availability-with-amazon-rds-proxy/>

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

D. Set up an Amazon RDS proxy for the database. Update the application to use the proxy endpoint.

upvoted 3 times

hachiri 1 year, 9 months ago

point is Aurora Multi-Master
Set up a secondary Aurora PostgreSQL cluster in the *same* AWS Region
upvoted 2 times

hachiri 1 year, 9 months ago

I mean correct is B
upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

Availability is the main requirement here. Even if RDS proxy is used, it will still find the writer instance unavailable during the scaling exercise.
Best option is to create an Amazon ElastiCache for Memcached cluster to handle the load during the scaling operation.
upvoted 1 times

pentium75 1 year, 5 months ago

"RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL "
<https://aws.amazon.com/de/blogs/database/improving-application-availability-with-amazon-rds-proxy/>
upvoted 2 times

wizcloudifa 1 year, 1 month ago

that will only account for the read operation part of it, we are concerned with the write operations here
upvoted 2 times

pentium75 1 year, 5 months ago

Failover is faster with RDS proxy
upvoted 2 times

AshishRocks 2 years ago

Set up an Amazon RDS proxy for the database. Update the application to use the proxy endpoint.
D is the answer
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #527

Topic 1

A company has a regional subscription-based streaming service that runs in a single AWS Region. The architecture consists of web servers and application servers on Amazon EC2 instances. The EC2 instances are in Auto Scaling groups behind Elastic Load Balancers. The architecture includes an Amazon Aurora global database cluster that extends across multiple Availability Zones.


The company wants to expand globally and to ensure that its application has minimal downtime.

Which solution will provide the MOST fault tolerance?

- A. Extend the Auto Scaling groups for the web tier and the application tier to deploy instances in Availability Zones in a second Region. Use an Aurora global database to deploy the database in the primary Region and the second Region. Use Amazon Route 53 health checks with a failover routing policy to the second Region.
- B. Deploy the web tier and the application tier to a second Region. Add an Aurora PostgreSQL cross-Region Aurora Replica in the second Region. Use Amazon Route 53 health checks with a failover routing policy to the second Region. Promote the secondary to primary as needed.
- C. Deploy the web tier and the application tier to a second Region. Create an Aurora PostgreSQL database in the second Region. Use AWS Database Migration Service (AWS DMS) to replicate the primary database to the second Region. Use Amazon Route 53 health checks with a failover routing policy to the second Region.
- D. Deploy the web tier and the application tier to a second Region. Use an Amazon Aurora global database to deploy the database in the primary Region and the second Region. Use Amazon Route 53 health checks with a failover routing policy to the second Region. Promote the secondary to primary as needed. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

TariqKipkemei **Highly Voted** 1 year, 10 months ago

Selected Answer: D

Auto Scaling groups can span Availability Zones, but not AWS regions.
Hence the best option is to deploy the web tier and the application tier to a second Region. Use an Amazon Aurora global

hence the best option is to deploy the web tier and the application tier to a second region. Use an Amazon Aurora global database to deploy the database in the primary Region and the second Region. Use Amazon Route 53 health checks with a failover routing policy to the second Region. Promote the secondary to primary as needed.

upvoted 21 times

awsgeek75 Most Recent 1 year, 5 months ago

Selected Answer: D

A: Not possible for autoscaling across regions

BC: Using PostgreSQL, not sure why?

D: MOST fault tolerant != MOST scalable. This gives least downtime.

upvoted 4 times

potomac 1 year, 7 months ago

Selected Answer: D

EC2 Auto Scaling groups are regional constructs. They can span Availability Zones, but not AWS regions

upvoted 3 times

thanhnv142 1 year, 7 months ago

527:

D is correct:

- B & C is not correct because it mentions Aurora PostgreSQL which is not mentioned in the question

- A is not correct because Auto scaling group can not span regions

upvoted 4 times

wsdasdasdqwda 1 year, 7 months ago

Simple as that.

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

Using an Aurora global database that spans both the primary and secondary regions provides automatic replication and failover capabilities for the database tier.

Deploying the web and application tiers to a second region provides fault tolerance for those components.

Using Route53 health checks and failover routing will route traffic to the secondary region if the primary region becomes unavailable.

This provides fault tolerance across all tiers of the architecture while minimizing downtime. Promoting the secondary database to primary ensures the second region can continue operating if needed.

A is close, but doesn't provide an automatic database failover capability.

B and C provide database replication, but not automatic failover.

So D is the most comprehensive and fault tolerant architecture.

upvoted 4 times

JA2018 6 months, 2 weeks ago

PostgreSQL in options b & C could be a distractor.... PostgreSQL not mentioned in stem

upvoted 1 times

Zox42 1 year, 11 months ago

Selected Answer: D

Answer D

upvoted 2 times

Zuit 1 year, 11 months ago

Selected Answer: D

D seems fitting: Global Database and deploying it in the new region

upvoted 2 times

MrAWSAssociate 1 year, 11 months ago

Selected Answer: B

B is correct!

upvoted 1 times

upvoted 1 times

manuh 1 year, 11 months ago

Replicated db doesnt mean they will act as a single db once the transfer is completed. Global db is the correct approach

upvoted 2 times

JA2018 6 months, 2 weeks ago

PostgreSQL in options b & C could be a distractor.... PostgreSQL not mentioned in stem

upvoted 1 times

r3mo 1 year, 12 months ago

"D" is the answer: because Aws Aurora Global Database allows you to read and write from any region in the global cluster. This enables you to distribute read and write workloads globally, improving performance and reducing latency. Data is replicated synchronously across regions, ensuring strong consistency.

upvoted 4 times

Henrytml 1 year, 12 months ago

Selected Answer: A

A is the only answer remain using ELB, both Web/App/DB has been taking care with replicating in 2nd region, lastly route 53 for failover over multiple regions

upvoted 1 times

Henrytml 1 year, 11 months ago

i will revoke my answer to standby web in 2nd region, instead of trigger to scale out

upvoted 1 times

manuh 1 year, 11 months ago

also Asg cant span beyond a region

upvoted 2 times

alexandercamachop 2 years ago

Selected Answer: D

B&C are discarded.

The answer is between A and D.

I would go with D because it explicitely created this web / app tier in second region, instead A just autoscales into a secondary region, rather then always having resources in this second region.

upvoted 4 times

JA2018 6 months, 2 weeks ago

as mentioned ASG cannot span across different AWS regions

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #528


Topic 1

A data analytics company wants to migrate its batch processing system to AWS. The company receives thousands of small data files periodically during the day through FTP. An on-premises batch job processes the data files overnight. However, the batch job takes hours to finish running.

The company wants the AWS solution to process incoming data files as soon as possible with minimal changes to the FTP clients that send the files. The solution must delete the incoming data files after the files have been processed successfully. Processing for each file needs to take 3-8 minutes.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. Use an Amazon EC2 instance that runs an FTP server to store incoming files as objects in Amazon S3 Glacier Flexible Retrieval. Configure a job queue in AWS Batch. Use Amazon EventBridge rules to invoke the job to process the objects nightly from S3 Glacier Flexible Retrieval. Delete the objects after the job has processed the objects.
- B. Use an Amazon EC2 instance that runs an FTP server to store incoming files on an Amazon Elastic Block Store (Amazon EBS) volume. Configure a job queue in AWS Batch. Use Amazon EventBridge rules to invoke the job to process the files nightly from the EBS volume. Delete the files after the job has processed the files.
- C. Use AWS Transfer Family to create an FTP server to store incoming files on an Amazon Elastic Block Store (Amazon EBS) volume. Configure a job queue in AWS Batch. Use an Amazon S3 event notification when each file arrives to invoke the job in AWS Batch. Delete the files after the job has processed the files.
- D. Use AWS Transfer Family to create an FTP server to store incoming files in Amazon S3 Standard. Create an AWS Lambda function to process the files and to delete the files after they are processed. Use an S3 event notification to invoke the Lambda function when the files arrive. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****pentium75** **Highly Voted** 11 months, 1 week ago**Selected Answer: D**

Obviously we choose AWS Transfer Family over hosting the FTP server ourselves on an EC2 instance. And "process incoming data files as soon as possible" -> trigger Lambda when files arrive. Lambda functions can run up to 15 minutes, it takes "3-8 minutes" per file -> works.

AWS Batch just schedules jobs, but these still need to run somewhere (Lambda, Fargate, EC2).

upvoted 5 times

wizcloudifa 7 months, 1 week ago

well if you opt for unmanaged compute option in AWS Batch in that case you dont have to worry about the compute, AWS batch automatically provisions it for you

upvoted 2 times

wsdasdasdqwda Most Recent 1 year, 1 month ago

FTP => AWS Transfer Family, => C or D, but in C is used EBS not S3 which needs EC2 and in general is more complex => very clear D.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

The key points:

Use AWS Transfer Family for the FTP server to receive files directly into S3. This avoids managing FTP servers.

Process each file as soon as it arrives using Lambda triggered by S3 events. Lambda provides fast processing time per file.

Lambda can also delete files after processing succeeds.

Options A, B, C involve more operational overhead of managing FTP servers and batch jobs. Processing latency would be higher waiting for batch windows.

Storing files in Glacier (Option A) adds latency for retrieving files.

upvoted 2 times

hsinchang 1 year, 4 months ago

Selected Answer: D

Processing for each file needs to take 3-8 minutes clearly indicates Lambda functions.

upvoted 2 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: D

Process incoming data files with minimal changes to the FTP clients that send the files = AWS Transfer Family.

Process incoming data files as soon as possible = S3 event notification.

Processing for each file needs to take 3-8 minutes = AWS Lambda function.

Delete file after processing = AWS Lambda function.

upvoted 4 times

antropaws 1 year, 5 months ago

Selected Answer: D

Most likely D.

upvoted 2 times

r3mo 1 year, 5 months ago

"D" Since each file takes 3-8 minutes to process the lambda function can process the data file without a problem.

upvoted 2 times

maver144 1 year, 5 months ago

Selected Answer: D

You cannot setup AWS Transfer Family to save files into EBS.

upvoted 3 times

oras2023 1 year, 5 months ago

<https://aws.amazon.com/aws-transfer-family/>

upvoted 2 times

secdgs 1 year, 5 months ago

Selected Answer: D

D. Because

1. process immediate when file transfer to S3 not wait for process several file in one time.
2. takes 3-8 can use Lambda.

C. Wrong because AWS Batch is use for run large-scale or large amount of data in one time.

upvoted 2 times

Aymanovitchy 1 year, 6 months ago

To meet the requirements of processing incoming data files as soon as possible with minimal changes to the FTP clients, and deleting the files after successful processing, the most operationally efficient solution would be:

D. Use AWS Transfer Family to create an FTP server to store incoming files in Amazon S3 Standard. Create an AWS Lambda function to process the files and delete them after processing. Use an S3 event notification to invoke the Lambda function when the files arrive.

upvoted 2 times

bajwa360 1 year, 6 months ago

Selected Answer: D

It should be D as lambda is more operationally viable solution given the fact each processing takes 3-8 minutes that lambda can handle

upvoted 2 times

alexandercamachop 1 year, 6 months ago

Selected Answer: C

Answer has to be between C or D.

Because Transfer Family is obvious do to FTP.

Now i would go with C because it uses AWS Batch, which makes more sense for Batch processing rather than AWS Lambda.

upvoted 1 times

pentium75 11 months, 1 week ago

Also Batch just triggers jobs, they still need to run somewhere (like in Lambda).

upvoted 2 times

pentium75 11 months, 1 week ago

Why? "Process incoming data files as soon as possible", by triggering the Lambda function when files arrive. Batch is for scheduled jobs.

upvoted 2 times

Bill1000 1 year, 6 months ago

I am between C and D. My reason is:

"The company wants the AWS solution to process incoming data files **as soon as possible** with minimal changes to the FTP clients that send the files."

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #529

Topic 1

A company is migrating its workloads to AWS. The company has transactional and sensitive data in its databases. The company wants to use AWS Cloud solutions to increase security and reduce operational overhead for the databases.

Which solution will meet these requirements?

- A. Migrate the databases to Amazon EC2. Use an AWS Key Management Service (AWS KMS) AWS managed key for encryption.
- B. Migrate the databases to Amazon RDS Configure encryption at rest. **Most Voted**
- C. Migrate the data to Amazon S3 Use Amazon Macie for data security and protection
- D. Migrate the database to Amazon RDS. Use Amazon CloudWatch Logs for data security and protection.

Correct Answer: B

Community vote distribution

B (100%)

Comments

AshishRocks **Highly Voted** 2 years ago

B is the answer

Why not C - Option C suggests migrating the data to Amazon S3 and using Amazon Macie for data security and protection. While Amazon Macie provides advanced security features for data in S3, it may not be directly applicable or optimized for databases, especially for transactional and sensitive data. Amazon RDS provides a more suitable environment for managing databases.

upvoted 11 times

oras2023 **Highly Voted** 2 years ago

Selected Answer: B

B. Migrate the databases to Amazon RDS Configure encryption at rest.
Looks like best option

upvoted 5 times

venomblade **Most Recent** 6 months, 3 weeks ago

Selected Answer: B

S3 and Macie works too. However, the question mentioned about transactional data. RDS is much better for running

transactional workloads. Thus, RDS is the correct answer here.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

A: Operational overhead of EC2 and whatever DB is running on it

C: Macie is not for data security, it's for identifying PII and sensitive data

D: CloudWatch is for cloud events and does not secure databases

B: RDS is managed so least operational overhead. Encryption at rest means security

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

Migrate the databases to Amazon RDS Configure encryption at rest.

upvoted 4 times

wendaz 1 year, 8 months ago

down voted.

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: B

Reduce Ops = Migrate the databases to Amazon RDS Configure encryption at rest

upvoted 4 times

alexandercamachop 2 years ago

Selected Answer: B

B for sure.

First the correct is Amazon RDS, then encryption at rest makes the database secure.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #530

Topic 1

A company has an online gaming application that has TCP and UDP multiplayer gaming capabilities. The company uses Amazon Route 53 to point the application traffic to multiple Network Load Balancers (NLBs) in different AWS Regions. The company needs to improve application performance and decrease latency for the online game in preparation for user growth.

Which solution will meet these requirements?

- A. Add an Amazon CloudFront distribution in front of the NLBs. Increase the Cache-Control max-age parameter.
- B. Replace the NLBs with Application Load Balancers (ALBs). Configure Route 53 to use latency-based routing.
- C. Add AWS Global Accelerator in front of the NLBs. Configure a Global Accelerator endpoint to use the correct listener ports. **Most Voted**
- D. Add an Amazon API Gateway endpoint behind the NLBs. Enable API caching. Override method caching for the different stages.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: C

The key considerations are:

The application uses TCP and UDP for multiplayer gaming, so Network Load Balancers (NLBs) are appropriate. AWS Global Accelerator can be added in front of the NLBs to improve performance and reduce latency by intelligently routing traffic across AWS Regions and Availability Zones.

Global Accelerator provides static anycast IP addresses that act as a fixed entry point to application endpoints in the optimal AWS location. This improves availability and reduces latency.

The Global Accelerator endpoint can be configured with the correct NLB listener ports for TCP and UDP.

upvoted 6 times

TariqKipkemei **Highly Voted** 1 year, 4 months ago

Selected Answer: C

TCP,UDP, Gaming = global accelerator and Network Load Balancer

upvoted 5 times

awsgeek75 Most Recent 11 months ago

Selected Answer: C

A: CloudFront is for caching. Not required

B: ALB is for HTTP layer, won't help with TCP UDP issues

D: API Gateway, API Caching total rubbish, ignore this option

C: Is correct as Global Accelerator uses unicast for reducing latency globally.

upvoted 2 times

Henrytml 1 year, 5 months ago

Selected Answer: C

only b and c handle TCP/UDP, and C comes with accelerator to enhance performance

upvoted 2 times

manuh 1 year, 5 months ago

Does alb handle udp? Can u share a source?

upvoted 1 times

alexandercamachop 1 year, 6 months ago

Selected Answer: C

UDP and TCP is AWS Global accelerator as it works in the Transportation layer.

Now this with NLB is perfect.

upvoted 3 times

oras2023 1 year, 6 months ago

Selected Answer: C

C is helping to reduce latency for end clients

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #531

Topic 1

A company needs to integrate with a third-party data feed. The data feed sends a webhook to notify an external service when new data is ready for consumption. A developer wrote an AWS Lambda function to retrieve data when the company receives a webhook callback. The developer must make the Lambda function available for the third party to call.

Which solution will meet these requirements with the MOST operational efficiency?

A. Create a function URL for the Lambda function. Provide the Lambda function URL to the third party for the webhook.

Most Voted

B. Deploy an Application Load Balancer (ALB) in front of the Lambda function. Provide the ALB URL to the third party for the webhook.

C. Create an Amazon Simple Notification Service (Amazon SNS) topic. Attach the topic to the Lambda function. Provide the public hostname of the SNS topic to the third party for the webhook.

D. Create an Amazon Simple Queue Service (Amazon SQS) queue. Attach the queue to the Lambda function. Provide the public hostname of the SQS queue to the third party for the webhook.

Correct Answer: A

Community vote distribution

A (100%)

Comments

TariqKipkemei Highly Voted 1 year, 10 months ago

Selected Answer: A

A function URL is a dedicated HTTP(S) endpoint for your Lambda function. When you create a function URL, Lambda automatically generates a unique URL endpoint for you.

upvoted 9 times

james2033 Highly Voted 1 year, 10 months ago

Selected Answer: A

Keyword "Lambda function" and "webhook". See <https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-saas-furls.html#create-stripe-cfn-stack>

upvoted 6 times

Most Recent 11 months ago

emakia most recent 11 months, 2 weeks ago

Selected Answer: A

AWS Lambda can provide a URL to call using Function URLs. This is a relatively new feature in AWS Lambda that allows you to create HTTPS endpoints for your Lambda functions, making it easy to invoke the function directly over the web.

Key Features of Lambda Function URLs:

Direct Access: Provides a simple and direct way to call a Lambda function via an HTTP(S) request.

Easy Configuration: You can create a function URL for a Lambda function using the AWS Management Console, AWS CLI, or AWS SDKs.

Managed Service: AWS manages the infrastructure for you, handling scaling, patching, and maintenance.

Security: You can configure authentication and authorization using AWS IAM or AWS Lambda function URL settings.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

Apart from simplest and most operational, I think A is the only option that will work!

BCD cannot even be implemented in real world imho. Happy to be corrected

upvoted 3 times

Orit 1 year, 6 months ago

B is the answerThe best solution to make the Lambda function available for the third party to call with the MOST operational efficiency is to deploy an Application Load Balancer (ALB) in front of the Lambda function and provide the ALB URL to the third party for the webhook. This solution is the most efficient because it allows the third party to call the Lambda function without having to worry about managing the Lambda function's availability or scaling. The ALB will automatically distribute traffic across multiple Lambda functions, if necessary, and will also provide redundancy in case of a failure.

upvoted 1 times

hro 1 year, 2 months ago

I believe you are correct. Lambda functions as targets - implementing ALBs

<https://docs.aws.amazon.com/elasticloadbalancing/latest/application/lambda-functions.html>

upvoted 1 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

The key points:

A Lambda function needs to be invoked by a third party via a webhook.

Using a function URL provides a direct invoke endpoint for the Lambda function. This is simple and efficient.

Options B, C, and D insert unnecessary components like ALB, SNS, SQS between the webhook and the Lambda function. These add complexity without benefit.

A function URL can be generated and provided to the third party quickly without additional infrastructure.

upvoted 5 times

Abrar2022 1 year, 11 months ago

Selected Answer: A

key word: Lambda function URLs

upvoted 3 times

maver144 1 year, 12 months ago

Selected Answer: A

<https://docs.aws.amazon.com/lambda/latest/dg/lambda-urls.html>

upvoted 2 times

jkhan2405 1 year, 12 months ago

Selected Answer: A

It's A

upvoted 2 times

alexandercamachop 2 years ago

Selected Answer: A

A would seem like the correct one but not sure.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #532

Topic 1

A company has a workload in an AWS Region. Customers connect to and access the workload by using an Amazon API Gateway REST API. The company uses Amazon Route 53 as its DNS provider. The company wants to provide individual and secure URLs for all customers.

Which combination of steps will meet these requirements with the MOST operational efficiency? (Choose three.)

A. Register the required domain in a registrar. Create a wildcard custom domain name in a Route 53 hosted zone and record in the zone that points to the API Gateway endpoint. **Most Voted**

B. Request a wildcard certificate that matches the domains in AWS Certificate Manager (ACM) in a different Region.

C. Create hosted zones for each customer as required in Route 53. Create zone records that point to the API Gateway endpoint.


D. Request a wildcard certificate that matches the custom domain name in AWS Certificate Manager (ACM) in the same Region. **Most Voted**

E. Create multiple API endpoints for each customer in API Gateway.

F. Create a custom domain name in API Gateway for the REST API. Import the certificate from AWS Certificate Manager (ACM). **Most Voted**

Correct Answer: ADF

Community vote distribution



Comments

AshishRocks **Highly Voted** 1 year, 11 months ago

Step A involves registering the required domain in a registrar and creating a wildcard custom domain name in a Route 53 hosted zone. This allows you to map individual and secure URLs for all customers to your API Gateway endpoints.

Step D is to request a wildcard certificate from AWS Certificate Manager (ACM) that matches the custom domain name you created in Step A. This wildcard certificate will cover all subdomains and ensure secure HTTPS communication.

Step F is to create a custom domain name in API Gateway for your REST API. This allows you to associate the custom domain name with your API Gateway endpoints and import the certificate from ACM for secure communication.

upvoted 9 times

Guru4Cloud Highly Voted 1 year, 9 months ago

Selected Answer: ADF

The key points:

Using a wildcard domain and certificate avoids managing individual domains/certs per customer. This is more efficient. The domain, hosted zone, and certificate should all be in the same region as the API Gateway REST API for simplicity. Creating multiple API endpoints per customer (Option E) adds complexity and is not required. Option B and C add unnecessary complexity by separating domains, certificates, and hosted zones.

upvoted 8 times

awsgeek75 Most Recent 1 year, 4 months ago

ADF looks right but not sure why C is wrong:

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-to-api-gateway.html#routing-to-api-gateway-config>

upvoted 3 times

LeonSauveterre 6 months ago

My guess on why C is not ideal: Creating individual hosted zones for each customer is too much. It definitely increases operational overhead, while using a wildcard domain eliminates this need (which is more ideal).

upvoted 1 times

ukivanlampi 1 year, 10 months ago

Selected Answer: ADF

<https://docs.aws.amazon.com/apigateway/latest/developerguide/how-to-custom-domains.html>

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/AboutHZWorkingWith.html>

upvoted 5 times

jaydesai8 1 year, 11 months ago

Selected Answer: ADF

ADF - makes sense

upvoted 3 times

jkhan2405 1 year, 12 months ago

Selected Answer: ADF

It's ADF

upvoted 3 times

MAMADOUG 1 year, 12 months ago

For me AFD

upvoted 3 times

alexandercamachop 2 years ago

Selected Answer: ADF

ADF - One to create the custom domain in Route 53 (Amazon DNS)

Second to request wildcard certificate from ADM

Thirds to import the certificate from ACM.

upvoted 3 times

AncaZalog 2 years ago

is ADF

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #533

Topic 1


A company stores data in Amazon S3. According to regulations, the data must not contain personally identifiable information (PII). The company recently discovered that S3 buckets have some objects that contain PII. The company needs to automatically detect PII in S3 buckets and to notify the company's security team.

Which solution will meet these requirements?

- A. Use Amazon Macie. Create an Amazon EventBridge rule to filter the SensitiveData event type from Macie findings and to send an Amazon Simple Notification Service (Amazon SNS) notification to the security team. **Most Voted**
- B. Use Amazon GuardDuty. Create an Amazon EventBridge rule to filter the CRITICAL event type from GuardDuty findings and to send an Amazon Simple Notification Service (Amazon SNS) notification to the security team.
- C. Use Amazon Macie. Create an Amazon EventBridge rule to filter the SensitiveData:S3Object/Personal event type from Macie findings and to send an Amazon Simple Queue Service (Amazon SQS) notification to the security team.
- D. Use Amazon GuardDuty. Create an Amazon EventBridge rule to filter the CRITICAL event type from GuardDuty findings and to send an Amazon Simple Queue Service (Amazon SQS) notification to the security team.

Correct Answer: A

Community vote distribution



Comments

alexandercamachop **Highly Voted** 2 years ago

Selected Answer: A

B and D are discarded as Macie is to identify PII.
Now that we have between A and C.

SNS is more suitable for this option as a pub/sub service, we subscribe the security team and then they will receive the notifications.

upvoted 15 times

LeonSauveterre **Most Recent** 6 months ago

Selected Answer: A

Macie + SNS = Identify PII + Notification

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

BD: Wrong products

AC: Uses Macie which is the right product but C uses SQS to notify security team which is an incomplete solution (what's listening to SQS?)

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: A

Detect PII -> Macie, A or C

Notify security team -> SNS, A or B

upvoted 4 times

potomac 1 year, 7 months ago

Selected Answer: A

C is SQS, not SNS

upvoted 4 times

Wayne23Fang 1 year, 9 months ago

SQS mentioned in C.

upvoted 2 times

Ale1973 1 year, 10 months ago

Selected Answer: A

Amazon SQS is typically used for decoupling and managing messages between distributed application components. It's not typically used for sending notifications directly to humans. On my opinion C isn't a best practice

upvoted 2 times

Kp88 1 year, 10 months ago

Those who say C , please read carefully (I made the same mistake lol). Teams can't be notified with SQS hence A.

upvoted 3 times

ukivanlamipi 1 year, 10 months ago

Selected Answer: C

there are different type of sensitive data: <https://docs.aws.amazon.com/macie/latest/user/findings-types.html>. if the question only focus on PII, then C is the answer. however, in reality, you will use A, because you will not want bank card, credential...etc all sensitive data , not only PII

upvoted 3 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: A

Automatically detect PII in S3 buckets = Amazon Macie

Notify security team = Amazon SNS

Trigger notification based on SensitiveData event type from Macie findings = EventBridge

upvoted 2 times

NASHDBA 1 year, 11 months ago

Selected Answer: C

There are different types of Sensitive Data. Here we are only referring to PII. Hence SensitiveData:S3Object/Personal. to use SNS, the security team must subscribe. SQS sends the information as designed

upvoted 1 times

narddrer 1 year, 11 months ago

Selected Answer: C

SensitiveData:S3Object/Personal

upvoted 1 times

randy1978 1 year, 11 months ago

jaydesai8 1 year, 11 months ago

Selected Answer: A

Sensitive = MACIE, and SNS to sent notification to the Security Team
upvoted 3 times

Iragmt 1 year, 11 months ago

C. Because the question mentioned PII only, there are other Sensitive Data aside from PII.
reference: <https://docs.aws.amazon.com/macie/latest/user/findings-publish-event-schemas.html> look for Event example for a sensitive data finding
upvoted 2 times

Ale1973 1 year, 10 months ago

But Amazon SQS is typically used for decoupling and managing messages between distributed application components. It's not typically used for sending notifications directly to humans!
upvoted 3 times

kapit 1 year, 11 months ago

AAAAAAA
upvoted 2 times

jack79 1 year, 11 months ago

C <https://docs.aws.amazon.com/macie/latest/user/findings-types.html>
and notice the ensitiveData:S3Object/Personal
The object contains personally identifiable information (such as mailing addresses or driver's license identification numbers), personal health information (such as health insurance or medical identification numbers), or a combination of the two.
upvoted 3 times

Ale1973 1 year, 10 months ago

But Amazon SQS is typically used for decoupling and managing messages between distributed application components. It's not typically used for sending notifications directly to humans!
upvoted 2 times

MAMADOUG 1 year, 12 months ago

I vote for A, Sensitive = MACIE, and SNS to prevent Security Team
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #534

Topic 1


A company wants to build a logging solution for its multiple AWS accounts. The company currently stores the logs from all accounts in a centralized account. The company has created an Amazon S3 bucket in the centralized account to store the VPC flow logs and AWS CloudTrail logs. All logs must be highly available for 30 days for frequent analysis, retained for an additional 60 days for backup purposes, and deleted 90 days after creation.

Which solution will meet these requirements MOST cost-effectively?

- A. Transition objects to the S3 Standard storage class 30 days after creation. Write an expiration action that directs Amazon S3 to delete objects after 90 days.
- B. Transition objects to the S3 Standard-Infrequent Access (S3 Standard-IA) storage class 30 days after creation. Move all objects to the S3 Glacier Flexible Retrieval storage class after 90 days. Write an expiration action that directs Amazon S3 to delete objects after 90 days.
- C. Transition objects to the S3 Glacier Flexible Retrieval storage class 30 days after creation. Write an expiration action that directs Amazon S3 to delete objects after 90 days. **Most Voted**
- D. Transition objects to the S3 One Zone-Infrequent Access (S3 One Zone-IA) storage class 30 days after creation. Move all objects to the S3 Glacier Flexible Retrieval storage class after 90 days. Write an expiration action that directs Amazon S3 to delete objects after 90 days.

Correct Answer: C

Community vote distribution



Comments

alexandercamachop **Highly Voted** 2 years ago

Selected Answer: C

C seems the most suitable.
Is the lowest cost.
After 30 days is backup only, doesn't specify frequent access.
Therefore we must transition the items after 30 days to Glacier Flexible Retrieval.

Also it says deletion after 90 days, so all answers specifying a transition after 90 days makes no sense.
upvoted 18 times

MAMADOUG 1 year, 12 months ago

Agree with you
upvoted 3 times

deechean **Highly Voted** 1 year, 9 months ago

Selected Answer: A

The Glacier min storage duration is 90 days. All the options using Glacier are wrong. Only A is feasible.
upvoted 10 times

daniel33 1 year, 8 months ago

S3 Standard is priced at \$0.023 per GB for the first 50 TB stored per month
S3 Glacier Flexible Retrieval costs \$0.0036 per GB stored per month
If you move or delete data in Glacier within 90-days since their creation, you will pay an additional charge, that is called an early deletion fee. In US East you will pay \$0.004/GB if you have deleted 1 GB in 2 months, \$0.008/GB if you have deleted 1 GB in 1 month and \$0.012 if you have deleted 1 GB within 3 months.

Even with the early deletion fee, it appears to me that answer 'A' would still be cheaper.
upvoted 3 times

awsgeek75 1 year, 4 months ago

But the objects are deleted after 90 days so how is this charge applicable?
upvoted 2 times

pentium75 1 year, 5 months ago

But why 'transition to the S3 Standard storage class', aren't they there already by default?
upvoted 4 times

LeonSauveterre **Most Recent** 6 months ago

Selected Answer: C

A: When logs are created, they are initially automatically created to S3 Standard. What are you transitioning anyway?
B & D: You could've just deleted logs after 90 days. No need to move them somewhere else then delete them. That's just money down the drain.
upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

C: Lowest cost
upvoted 2 times

awsgeek75 1 year, 5 months ago

A: Standard storage is default so this is wrong.
B: Looks wrong because it moves object to S3GFR after 90 days when they could just be deleted so extra cost
D: Same problem as B
upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: C

Not A: Objects are created in S3 Standard, so it doesn't make sense to 'transition' them there "30 days after creation"
Not B or C: No need to "move all objects to the S3 Glacier Flexible Retrieval storage class after 90 days" because we want to delete, not archive, them. Even if we would delete them right after moving, we would pay 90 days minimum storage duration. Plus, we are using "Infrequent Access" classes here, but we have no access at all.
upvoted 3 times

Rhydian25 11 months, 2 weeks ago

I guess you wanted to write "Not B or D"
upvoted 3 times

ftaws 1 year, 5 months ago

Selected Answer: A

requirement : frequently analysis
search cost : S3 STD 0.0004 vs IA 0.001
so IA is more expensive than STD(A)
upvoted 1 times

EdenWang 1 year, 6 months ago

Selected Answer: C

C is most cost-effective
upvoted 3 times

Hades2231 1 year, 9 months ago

Selected Answer: C

Things to note are: 30 days frequent access and 90 days after creation, so you only need to do 2 things, not 3. Objects in S3 will be stored by default for 30 days before you can move it to somewhere else, so C is the answer.
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>
upvoted 5 times

rjbihari 1 year, 9 months ago

C is the correct one .
As after 30 days it doesn't says about access / retrieval , only backup so move items after 30 days to Glacier Flexible Retrieval. And after it says deletion , so expiration action will ensure that the objects are deleted after 90 days, even if they are not accessed
upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

I think - it is B
The first 30 days, the logs need to be highly available for frequent analysis. The S3 Standard storage class is the most expensive storage class, but it also provides the highest availability.
After 30 days, the logs still need to be retained for backup purposes, but they do not need to be accessed frequently. The S3 Standard-IA storage class is a good option for this, as it is less expensive than the S3 Standard storage class.
After 90 days, the logs can be moved to the S3 Glacier Flexible Retrieval storage class. This is the most cost-effective storage class for long-term archiving.
The expiration action will ensure that the objects are deleted after 90 days, even if they are not accessed
upvoted 2 times

pentium75 1 year, 5 months ago

"After 90 days, the logs can be moved to the S3 Glacier Flexible Retrieval storage class. This is the most cost-effective storage class for long-term archiving." yeah but we don't need long-term archiving, we want to delete them after 90 days.
upvoted 3 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

C is the most cost effective solution.
upvoted 2 times

antropaws 1 year, 11 months ago

Selected Answer: C

C most likely.
upvoted 2 times

y0eri 1 year, 11 months ago

Selected Answer: A

Question says "All logs must be highly available for 30 days for frequent analysis" I think the answer is A. Glacier is not made for frequent access.
upvoted 2 times

y0eri 1 year, 11 months ago

I take that back. Moderator, please delete my comment.
upvoted 4 times

KMohsoe 1 year, 12 months ago

Selected Answer: B

I think B
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #535

Topic 1

A company is building an Amazon Elastic Kubernetes Service (Amazon EKS) cluster for its workloads. All secrets that are stored in Amazon EKS must be encrypted in the Kubernetes etcd key-value store.

Which solution will meet these requirements?

- A. Create a new AWS Key Management Service (AWS KMS) key. Use AWS Secrets Manager to manage, rotate, and store all secrets in Amazon EKS.
- B. Create a new AWS Key Management Service (AWS KMS) key. Enable Amazon EKS KMS secrets encryption on the Amazon EKS cluster. **Most Voted**
- C. Create the Amazon EKS cluster with default options. Use the Amazon Elastic Block Store (Amazon EBS) Container Storage Interface (CSI) driver as an add-on.
- D. Create a new AWS Key Management Service (AWS KMS) key with the alias/aws/ebs alias. Enable default Amazon Elastic Block Store (Amazon EBS) volume encryption for the account.

Correct Answer: B*Community vote distribution***Comments****Guru4Cloud** **Highly Voted** 9 months, 2 weeks ago**Selected Answer: B**

B is the correct solution to meet the requirement of encrypting secrets in the etcd store for an Amazon EKS cluster.

The key points:

Create a new KMS key to use for encryption.

Enable EKS secrets encryption using that KMS key on the EKS cluster. This will encrypt secrets in the Kubernetes etcd store.

Option A uses Secrets Manager which does not encrypt the etcd store.

Option C uses EBS CSI which is unrelated to etcd encryption.

Option D enables EBS encryption but does not address etcd encryption.

upvoted 7 times

TariqKipkemei **Highly Voted** 10 months, 2 weeks ago**Selected Answer: B**

EKS supports using AWS KMS keys to provide envelope encryption of Kubernetes secrets stored in EKS. Envelope encryption adds an additional, customer-managed layer of encryption for application secrets or user data that is stored within a Kubernetes cluster.

<https://eksctl.io/usage/kms-encryption/>
upvoted 6 times

FlyingHawk Most Recent 4 months, 2 weeks ago

Selected Answer: B

The problem is to ensure all secrets in Amazon EKS be encrypted, so the correct solution is to enable the secret encryption on the EKS cluster with a new AWS Key:

```
eksctl utils enable-secrets-encryption \
```

```
--cluster my-cluster \
```

```
--key-arn arn:aws:kms:region-code:account:key/key
```

https://docs.aws.amazon.com/en_en/eks/latest/userguide/enable-kms.html

upvoted 1 times

LeonSauveterre 6 months ago

Selected Answer: B

"must be encrypted in the Kubernetes etcd key-value store" - that leaves B only.

upvoted 1 times

manuh 11 months, 2 weeks ago

Selected Answer: A

Why not a

upvoted 1 times

TariqKipkemei 10 months, 2 weeks ago

option A does not enable Amazon EKS KMS secrets encryption on the Amazon EKS cluster

upvoted 2 times

MrAWSAssociate 11 months, 3 weeks ago

Selected Answer: B

B is the right option.

<https://docs.aws.amazon.com/eks/latest/userguide/enable-kms.html>

upvoted 5 times

alexandercamachop 1 year ago

Selected Answer: B

It is B, because we need to encrypt inside of the EKS cluster, not outside.

AWS KMS is to encrypt at rest.

upvoted 5 times

AncaZalog 1 year ago

is B, not D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #536

Topic 1

A company wants to provide data scientists with near real-time read-only access to the company's production Amazon RDS for PostgreSQL database. The database is currently configured as a Single-AZ database. The data scientists use complex queries that will not affect the production database. The company needs a solution that is highly available.

Which solution will meet these requirements MOST cost-effectively?

- A. Scale the existing production database in a maintenance window to provide enough power for the data scientists.
- B. Change the setup from a Single-AZ to a Multi-AZ instance deployment with a larger secondary standby instance. Provide the data scientists access to the secondary instance.
- C. Change the setup from a Single-AZ to a Multi-AZ instance deployment. Provide two additional read replicas for the data scientists.
- D. Change the setup from a Single-AZ to a Multi-AZ cluster deployment with two readable standby instances. Provide read endpoints to the data scientists. **Most Voted**

Correct Answer: D*Community vote distribution*

D (82%)

Other (18%)

Comments**NASHDBA** **Highly Voted** 1 year, 11 months ago**Selected Answer: D**

Highly Available = Multi-AZ Cluster

Read-only + Near Real time = readable standby.

Read replicas are async whereas readable standby is synchronous.

<https://stackoverflow.com/questions/70663036/differences-b-w-aws-read-replica-and-the-standby-instances>

upvoted 23 times

chickenmf 1 year, 2 months ago

a Multi-AZ instance deployment is also highly available for a lower cost, no?

upvoted 2 times

Smart 1 year, 9 months ago

This^ is the reason.

upvoted 3 times

maver144 Highly Voted 1 year, 12 months ago

It's either C or D. To be honest, I find the newest questions to be ridiculously hard (roughly 500+). I agree with @alexandercamachop that Multi Az in Instance mode is cheaper than Cluster. However, with Cluster we have reader endpoint available to use out-of-box, so there is no need to provide read-replicas, which also has its own costs. The ridiculous part is that I'm pretty sure even the AWS support would have troubles to answer which configuration is MOST cost-effective.

upvoted 13 times

maver144 1 year, 12 months ago

Near real-time is clue for C, since read replicas are async, but still its not obvious question.

upvoted 2 times

manuh 1 year, 11 months ago

Absolutely true that 500+ questions are damn difficult to answer. I still dont know why is B incorrect. Shouldn't 1 extra be better than 2 ?

upvoted 1 times

cyber_bedouin 1 year, 6 months ago

they are not all hard, most are normal. its just this one and that one about EKS encryption control plane (earlier than this page).

upvoted 3 times

AdamVigas Most Recent 3 months, 3 weeks ago

Selected Answer: C

two readable standby instances...if they are standby they are not supposed to read...am i getting something wrong ? i would go for C

upvoted 2 times

emakid 11 months, 2 weeks ago

Selected Answer: C

Option D: Multi-AZ cluster deployment with two readable standby instances would be more costly and is not necessary if read replicas are sufficient for the data scientists' needs.

Thus, Option C is the most cost-effective and operationally efficient solution to meet the company's requirements.

upvoted 2 times

osmk 1 year, 3 months ago

Selected Answer: D

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/multi-az-db-clusters-concepts.html>

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: D

Not A - Not "highly available"

Not B - "Access to the secondary instance" is not possible in Multi-AZ

Not C - Multi-AZ + two (!) read replicas is more expensive than cluster

D - Provides "readable standby instances"

upvoted 3 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

D

<https://aws.amazon.com/about-aws/whats-new/2023/01/amazon-rds-multi-az-readable-standbys-rds-postgresql-inbound-replication/>

upvoted 3 times

bogobob 1 year, 6 months ago

Selected Answer: D

<https://aws.amazon.com/blogs/database/choose-the-right-amazon-rds-deployment-option-single-az-instance-multi-az-instance-or-multi-az-database-cluster/>

C would mean you are paying for 4 instances (primary, backup, and 2 read instances). D would be 3 (primary, and 2 backup). Difficult to be sure, pricing calculator doesn't even include clusters yet.

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

Option D is the most cost-effective solution that meets the requirements for this scenario.

The key considerations are:

Data scientists need read-only access to near real-time production data without affecting performance.

High availability is required.

Cost should be minimized.

upvoted 2 times

ukivanlamipi 1 year, 10 months ago

Selected Answer: D

<https://aws.amazon.com/blogs/database/choose-the-right-amazon-rds-deployment-option-single-az-instance-multi-az-instance-or-multi-az-database-cluster/>

only multi AZ cluster have reader endpoint. multi AZ instance secondary replicate is not allow to access

upvoted 2 times

msdnpro 1 year, 10 months ago

Selected Answer: D

Support for D:

Amazon RDS now offers Multi-AZ deployments with readable standby instances (also called Multi-AZ DB cluster deployments) in preview. You should consider using Multi-AZ DB cluster deployments with two readable DB instances if you need additional read capacity in your Amazon RDS Multi-AZ deployment and if your application workload has strict transaction latency requirements such as single-digit milliseconds transactions.

<https://aws.amazon.com/blogs/database/readable-standby-instances-in-amazon-rds-multi-az-deployments-a-new-high-availability-option/>

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: D

Unlike Multi-AZ instance deployment, where the secondary instance can't be accessed for read or writes, Multi-AZ DB cluster deployment consists of primary instance running in one AZ serving read-write traffic and two other standby running in two different AZs serving read traffic.

upvoted 4 times

Iragmt 1 year, 11 months ago

Selected Answer: D

D. using Multi-AZ DB cluster deployments with two readable DB instances if you need additional read capacity in your Amazon RDS Multi-AZ deployment and if your application workload has strict transaction latency requirements such as single-digit milliseconds transactions. <https://aws.amazon.com/blogs/database/readable-standby-instances-in-amazon-rds-multi-az-deployments-a-new-high-availability-option/>

while on read replicas, Amazon RDS then uses the asynchronous replication method for the DB engine to update the read replica whenever there is a change to the primary DB instance.

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_ReadRepl.html

upvoted 2 times

manuh 1 year, 11 months ago

Selected Answer: B

Why not b. Shouldnt it have less number of instances than both c and d?

upvoted 2 times

baba365 1 year, 11 months ago

Complex queries on single db will affect performance of db
upvoted 2 times

pentium75 1 year, 5 months ago

You can't 'access the secondary instance' as suggested by B
upvoted 2 times

pentium75 1 year, 5 months ago

"The Multi-AZ instance is suitable for business/mission critical applications that require high availability with low RTO/RPO and resilience to availability zone outage. However, this high availability option isn't a scaling solution for read-only scenarios. You can't use a standby replica to serve read traffic. To serve read-only traffic, use a Multi-AZ DB cluster or a read replica instead."
upvoted 2 times

baba365 1 year, 11 months ago

Multi-AZ is about twice the price of Single-AZ. For example:
db.t2.micro single - \$0.017/hour
db.t2.micro multi - \$0.034/hour

option C: 1 primary + 1 standby + 2 replica = 4Db
option D: 1 primary + 2 standby = 3Db

D. appears to be most cost effective
upvoted 3 times

wsdasdasdqwdaw 1 year, 7 months ago

I think the best explanation I've read so far.
upvoted 2 times

0628atv 1 year, 11 months ago

D:
<https://aws.amazon.com/tw/blogs/database/readable-standby-instances-in-amazon-rds-multi-az-deployments-a-new-high-availability-option/>
upvoted 2 times

vrevkov 1 year, 11 months ago

Selected Answer: D

Forgot to vote
upvoted 3 times

vrevkov 1 year, 11 months ago

I think it's D.
C: Multi-AZ instance = active + standby + two read replicas = 4 RDS instances
D: Multi-AZ cluster = Active + two standby = 3 RDS instances

Single-AZ and Multi-AZ deployments: Pricing is billed per DB instance-hour consumed from the time a DB instance is launched until it is stopped or deleted.
<https://aws.amazon.com/rds/postgresql/pricing/?pg=pr&loc=3>
In the case of a cluster, you will pay less.
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #537

Topic 1

A company runs a three-tier web application in the AWS Cloud that operates across three Availability Zones. The application architecture has an Application Load Balancer, an Amazon EC2 web server that hosts user session states, and a MySQL database that runs on an EC2 instance. The company expects sudden increases in application traffic. The company wants to be able to scale to meet future application capacity demands and to ensure high availability across all three Availability Zones.

Which solution will meet these requirements?

- A. Migrate the MySQL database to Amazon RDS for MySQL with a Multi-AZ DB cluster deployment. Use Amazon ElastiCache for Redis with high availability to store session data and to cache reads. Migrate the web server to an Auto Scaling group that is in three Availability Zones. **Most Voted**
- B. Migrate the MySQL database to Amazon RDS for MySQL with a Multi-AZ DB cluster deployment. Use Amazon ElastiCache for Memcached with high availability to store session data and to cache reads. Migrate the web server to an Auto Scaling group that is in three Availability Zones.
- C. Migrate the MySQL database to Amazon DynamoDB. Use DynamoDB Accelerator (DAX) to cache reads. Store the session data in DynamoDB. Migrate the web server to an Auto Scaling group that is in three Availability Zones.
- D. Migrate the MySQL database to Amazon RDS for MySQL in a single Availability Zone. Use Amazon ElastiCache for Redis with high availability to store session data and to cache reads. Migrate the web server to an Auto Scaling group that is in three Availability Zones.

Correct Answer: A*Community vote distribution*

A (80%)

B (20%)

Comments**alexandercamachop** **Highly Voted** 1 year, 6 months ago**Selected Answer: A**

Memcached is best suited for caching data, while Redis is better for storing data that needs to be persisted. If you need to store data that needs to be accessed frequently, such as user profiles, session data, and application settings, then Redis is the better choice

upvoted 16 times

nonameforyou 1 year, 5 months ago

nonameforyou 1 year, 5 months ago

and for high availability, it's better than memcached
upvoted 2 times

nonameforyou 1 year, 5 months ago

but does rds multi-az provide the needed scalability?
upvoted 3 times

wsdasdasdqwdaw 1 year, 1 month ago

it is multi-az cluster deployment, same as B, so yes, it is providing the needed scalability. Great explanation.
upvoted 2 times

AncaZalog Highly Voted 1 year, 6 months ago

is A not B
upvoted 5 times

zdi561 Most Recent 4 months ago

Selected Answer: A

AI answer, For managing user sessions on AWS ElastiCache, Redis is generally the better choice compared to Memcached because Redis offers more advanced data structures, built-in persistence features, and better support for complex session data, while Memcached is simpler but only suitable for basic key-value caching with limited data types

upvoted 1 times

osmk 9 months, 3 weeks ago

Selected Answer: A

Replication: Redis supports creating multiple replicas for read scalability and high availability.<https://aws.amazon.com/elasticache/redis-vs-memcached/>

upvoted 2 times

awsgeek75 10 months, 3 weeks ago

Selected Answer: A

A because of "Amazon EC2 web server that hosts user session states"

C: RDS to DynamoDB doesn't make total sense

D: Single zone is not HA

Between A and B, A is suitable because of session state and ElastiCache with Redis is more HA than option B

upvoted 2 times

mr123dd 11 months, 1 week ago

Selected Answer: A

B: from what I know, Memcached provide better performance and simplicity but lower availability than redis.

C: mysql is relational database, dynamodb is nosql

D: single AZ

upvoted 2 times

pentium75 11 months, 1 week ago

Selected Answer: A

ElastiCache for Redis supports HA, ElastiCache for Memcached does not:
<https://docs.aws.amazon.com/AmazonElastiCache/latest/mem-ug/SelectEngine.html>

C could in theory work, but session data is typically stored in ElastiCache, not in DynamoDB.

D is not HA.

upvoted 3 times

Cyberkayu 11 months, 3 weeks ago

Selected Answer: B

'hosts user session states' in question, thus redis

hosts user session states in question, thus Redis
upvoted 2 times

pentium75 11 months, 1 week ago

Right, but Redis is A
upvoted 2 times

potomac 1 year, 1 month ago

Selected Answer: A

Redis is a widely adopted in-memory data store for use as a database, cache, message broker, queue, session store, and leaderboard.

<https://aws.amazon.com/elasticache/redis/>

upvoted 5 times

thanhnv142 1 year, 1 month ago

B is correct.

We are left with 2 options: A and B. But it requires that the system be able to scale to meet future application capacity demands. Redis is very good. But its drawback is not scalable. That's why they implement memcached.

upvoted 1 times

ErnShm 1 year, 2 months ago

A
Redis as an in-memory data store with high availability and persistence is a popular choice among application developers to store and manage session data for internet-scale applications. Redis provides the sub-millisecond latency, scale, and resiliency required to manage session data such as user profiles, credentials, session state, and user-specific personalization.

upvoted 2 times

Gajendr 11 months, 2 weeks ago

Redis provides replication while memcached doesn't.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

The key reasons why option A is preferable:

RDS Multi-AZ provides high availability for MySQL by synchronously replicating data across AZs. Automatic failover handles AZ outages.

ElastiCache for Redis is better suited for session data caching than Memcached. Redis offers more advanced data structures and flexibility.

Auto scaling across 3 AZs provides high availability for the web tier

upvoted 2 times

ukivanlampi 1 year, 4 months ago

Selected Answer: B

The difference between Redis and Memcache is that Memcache supports multithread process to handle the increase of application traffic. <https://aws.amazon.com/elasticache/redis-vs-memcached/>

upvoted 2 times

pentium75 11 months, 1 week ago

ElastiCache for Memcached says "No" for "High Availability"

<https://docs.aws.amazon.com/AmazonElastiCache/latest/mem-ug/SelectEngine.html>

upvoted 2 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: B

This requirement wins for me: "be able to scale to meet future application capacity demands".

Memcached implements a multi-threaded architecture, it can make use of multiple processing cores. This means that you can handle more operations by scaling up compute capacity.

<https://aws.amazon.com/elasticache/redis-vs-memcached/#:~:text=by%20their%20rank.,Multithreaded%20architecture,-Since%20Memcached%20is>

upvoted 1 times

plndmns 1 year, 4 months ago

cache reads is memcached right?

upvoted 1 times

MrAWSAssociate 1 year, 5 months ago

Selected Answer: B

B is correct!

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #538

Topic 1

A global video streaming company uses Amazon CloudFront as a content distribution network (CDN). The company wants to roll out content in a phased manner across multiple countries. The company needs to ensure that viewers who are outside the countries to which the company rolls out content are not able to view the content.

Which solution will meet these requirements?

A. Add geographic restrictions to the content in CloudFront by using an allow list. Set up a custom error message.

Most Voted

B. Set up a new URL for restricted content. Authorize access by using a signed URL and cookies. Set up a custom error message.

C. Encrypt the data for the content that the company distributes. Set up a custom error message.

D. Create a new URL for restricted content. Set up a time-restricted access policy for signed URLs.

Correct Answer: A

Community vote distribution

A (100%)

Comments

potomac Highly Voted 1 year, 1 month ago

Selected Answer: A

The CloudFront geographic restrictions feature lets you control distribution of your content at the country level for all files that you're distributing with a given web distribution.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/georestrictions.html>

upvoted 5 times

awsgeek75 Highly Voted 11 months ago

Selected Answer: A

This question asks us to guess Netflix subscription model in 2 mins! lol!

BCD are impractical for geo restrictions as you cannot restrict URL by region and you cannot encrypt by geo region (country etc)

upvoted 5 times