

might be heavy for a short period of time, relieved by long periods of high idleness or no activity at all.

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-serverless-v2.how-it-works.html>

upvoted 3 times

ersin13 1 year, 10 months ago

A. Migrate the existing RDS for MySQL database to an Aurora Serverless v2 MySQL database cluster.

upvoted 2 times

mrsoa 1 year, 10 months ago

Selected Answer: B

B seems to be the correct answer, because if we have a predictable workload Aurora database seems to be most cost effective however if we have unpredictable workload aurora serverless seems to be more cost effective because our database will scale up and down

for more informations please read this article

<https://medium.com/trackit/aurora-or-aurora-serverless-v2-which-is-more-cost-effective-bcd12e172dcf>

upvoted 3 times

Chef_couincouin 1 year, 7 months ago

according to the link, i understand that Aurora Serverless is ideal for sudden peaks in database usage with moderate or minimal usage during other periods of the day. So Answer is A

upvoted 3 times

Smart 1 year, 9 months ago

True but due to autoscaling - it will be cheaper...check example#1 in the your link.

upvoted 1 times

Smart 1 year, 9 months ago

Correct Answer is A

upvoted 2 times

pentium75 1 year, 5 months ago

Provisioned RDS (as in B) is good for steady (not "predictable") workloads. In this case, the workload is predictable, but the prediction is that it will be used only 2 hours per week.

upvoted 3 times

pentium75 1 year, 5 months ago

Aurora Serverless is for "infrequent, intermittent, OR unpredictable workloads"

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #575

Topic 1

A company deploys its applications on Amazon Elastic Kubernetes Service (Amazon EKS) behind an Application Load Balancer in an AWS Region. The application needs to store data in a PostgreSQL database engine. The company wants the data in the database to be highly available. The company also needs increased capacity for read workloads.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Create an Amazon DynamoDB database table configured with global tables.
- B. Create an Amazon RDS database with Multi-AZ deployments.
- C. Create an Amazon RDS database with Multi-AZ DB cluster deployment. **Most Voted**
- D. Create an Amazon RDS database configured with cross-Region read replicas.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: C

RDS Multi-AZ DB cluster deployments provide high availability, automatic failover, and increased read capacity. A multi-AZ cluster automatically handles replicating data across AZs in a single region. This maintains operational efficiency as it is natively managed by RDS without needing external replication. DynamoDB global tables involve complex provisioning and requires app changes. RDS read replicas require manual setup and management of replication. RDS Multi-AZ clustering is purpose-built by AWS for HA PostgreSQL deployments and balancing read workloads.

upvoted 9 times

luiscc **Highly Voted** 1 year, 10 months ago

Selected Answer: C

DB cluster deployment can scale read workloads by adding read replicas. This provides increased capacity for read workloads without impacting the write workload.

upvoted 5 times

MatAlves **Most Recent** 8 months, 3 weeks ago

Selected Answer: C

Selected Answer: C

"A Multi-AZ DB cluster deployment is a semisynchronous, high availability deployment mode of Amazon RDS with two readable replica DB instances."

upvoted 2 times

upliftinghut 1 year, 4 months ago

Selected Answer: C

multi-AZ addresses both HA & increased read capacity with synchronous data replication between main DB & standby. Read replica is not enough because only increased read capacity not enabling HA, besides the data replication is async

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: C

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/multi-az-db-clusters-concepts.html>

"A Multi-AZ DB cluster deployment is a semisynchronous, high availability deployment mode of Amazon RDS with two readable standby DB instances"

A: DynamoDB is not Postgres

B: Although HA is achieve but it does not increase the read capacity as much as C without additional operational complexity

D: Cross region is not a requirement and won't solve the same region HA or read issues

upvoted 2 times

aws94 1 year, 5 months ago

Selected Answer: C

Multi-AZ DB Cluster Deployment = Aurora

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

Multi-AZ DB cluster deployments provides two readable DB instances if you need additional read capacity

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: C

C is correct

upvoted 2 times

avkya 1 year, 10 months ago

Selected Answer: C

Multi-AZ DB clusters provide high availability, increased capacity for read workloads, and lower write latency when compared to Multi-AZ DB instance deployments.

upvoted 2 times

mrsoa 1 year, 10 months ago

Selected Answer: C

CCCCCCCCCcCCcCcCCCCccccCc

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #576

Topic 1

A company is building a RESTful serverless web application on AWS by using Amazon API Gateway and AWS Lambda. The users of this web application will be geographically distributed, and the company wants to reduce the latency of API requests to these users.

Which type of endpoint should a solutions architect use to meet these requirements?

- A. Private endpoint
- B. Regional endpoint
- C. Interface VPC endpoint

D. Edge-optimized endpoint **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

mrsoa **Highly Voted** 1 year, 4 months ago

Selected Answer: D

The correct answer is D

API Gateway - Endpoint Types

- Edge-Optimized (default): For global clients
- Requests are routed through the CloudFront Edge locations (improves latency)
- The API Gateway still lives in only one region
- Regional:
 - For clients within the same region
 - Could manually combine with CloudFront (more control over the caching strategies and the distribution)
- Private:
 - Can only be accessed from your VPC using an interface VPC endpoint (ENI)
 - Use a resource policy to define access

upvoted 8 times

potomac **Highly Voted** 1 year, 1 month ago

Selected Answer: D

An edge-optimized API endpoint typically routes requests to the nearest CloudFront Point of Presence (POP), which could help in cases where your clients are geographically distributed. This is the default endpoint type for API Gateway REST APIs.

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-endpoint-types.html>
upvoted 5 times

awsgeek75 Most Recent 10 months, 3 weeks ago

Selected Answer: D

geographically distributed users + low latency = Edge optimized endpoint

upvoted 2 times

TariqKipkemei 1 year ago

Selected Answer: D

An edge-optimized API endpoint typically routes requests to the nearest CloudFront Point of Presence (POP), which could help in cases where your clients are geographically distributed. This is the default endpoint type for API Gateway REST APIs.

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-endpoint-types.html#:~:text=API%20endpoint%20typically-,routes,-requests%20to%20the>
upvoted 3 times

dilaaziz 1 year, 1 month ago

Selected Answer: D

<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-endpoint-types.html>
upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

Edge-optimized endpoint

upvoted 3 times

Josantru 1 year, 4 months ago

Correct D.

Edge-optimized API endpoints

An edge-optimized API endpoint is best for geographically distributed clients. API requests are routed to the nearest CloudFront Point of Presence (POP). This is the default endpoint type for API Gateway REST APIs.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #577

Topic 1

A company uses an Amazon CloudFront distribution to serve content pages for its website. The company needs to ensure that clients use a TLS certificate when accessing the company's website. The company wants to automate the creation and renewal of the TLS certificates.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Use a CloudFront security policy to create a certificate.
- B. Use a CloudFront origin access control (OAC) to create a certificate.
- C. Use AWS Certificate Manager (ACM) to create a certificate. Use DNS validation for the domain. **Most Voted**
- D. Use AWS Certificate Manager (ACM) to create a certificate. Use email validation for the domain.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Bmaster** **Highly Voted** 1 year, 4 months ago

C is correct.

"ACM provides managed renewal for your Amazon-issued SSL/TLS certificates. This means that ACM will either renew your certificates automatically (if you are using DNS validation), or it will send you email notices when expiration is approaching. These services are provided for both public and private ACM certificates."

<https://docs.aws.amazon.com/acm/latest/userguide/managed-renewal.html>

upvoted 10 times

Guru4Cloud **Highly Voted** 1 year, 3 months ago**Selected Answer: C**

The key reasons are:

AWS Certificate Manager (ACM) provides free public TLS/SSL certificates and handles certificate renewals automatically. Using DNS validation with ACM is operationally efficient since it automatically makes changes to Route 53 rather than requiring manual validation steps.

ACM integrates natively with CloudFront distributions for delivering HTTPS content. CloudFront security policies and origin access controls do not issue TLS certificates.

CloudFront security policies and origin access controls do not issue TLS certificates.
Email validation requires manual steps to approve the domain validation emails for each renewal.

upvoted 6 times

awsgeek75 Most Recent 11 months ago

Selected Answer: C

For me, C is the only realistic option as I don't think you can do AB without a lot of complexity. D just makes no sense.

upvoted 2 times

ibu007 1 year, 3 months ago

Selected Answer: C

Use AWS Certificate Manager (ACM) to create a certificate. Use DNS validation for the domain

upvoted 4 times

chen0305_099 1 year, 3 months ago

Selected Answer: C

C 似乎是正確的

upvoted 4 times

Kiki_Pass 1 year, 4 months ago

Selected Answer: C

"DNS Validation is preferred for automation purposes" -- Stephane's course on Udemy

upvoted 3 times

mrsoa 1 year, 4 months ago

Selected Answer: C

C seems to be correct

upvoted 2 times

nananashi 1 year, 4 months ago

I think the general product uses DNS rather than email to automate, is the given answer correct?

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #578

Topic 1

A company deployed a serverless application that uses Amazon DynamoDB as a database layer. The application has experienced a large increase in users. The company wants to improve database response time from milliseconds to microseconds and to cache requests to the database.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use DynamoDB Accelerator (DAX). **Most Voted**


B. Migrate the database to Amazon Redshift.

C. Migrate the database to Amazon RDS.

D. Use Amazon ElastiCache for Redis.

Correct Answer: A

Community vote distribution



Comments

h8er **Highly Voted** 1 year, 10 months ago

Selected Answer: A

Amazon DynamoDB Accelerator (DAX) is a fully managed, highly available, in-memory cache for Amazon DynamoDB that delivers up to a 10 times performance improvement—from milliseconds to microseconds—even at millions of requests per second.

[https://aws.amazon.com/dynamodb/dax/#:~:text=Amazon%20DynamoDB%20Accelerator%20\(DAX\)%20is,millions%20of%20re,quests%20per%20second.](https://aws.amazon.com/dynamodb/dax/#:~:text=Amazon%20DynamoDB%20Accelerator%20(DAX)%20is,millions%20of%20re,quests%20per%20second.)

upvoted 11 times

MatAlves **Most Recent** 8 months, 3 weeks ago

Amazon ElastiCache for Redis would help with "caching requests", but not "improve database response" itself.

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

DAX is least operations overhead.

B: Redshift, although powerful, but is for analytics

C: Downgrading to RDS won't help

D: EC for Redis is more for persistent caching so would be good but lot of operational overhead

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

improve DynamoDB response time from milliseconds to microseconds and to cache requests to the database = DynamoDB Accelerator (DAX)

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

Use DynamoDB Accelerator (DAX).

upvoted 2 times

pentium75 1 year, 5 months ago

Which is A, not C.

upvoted 3 times

awsgeek75 1 year, 4 months ago

Quote A but mark C. You need more coffee mate :)

upvoted 5 times

mrsoa 1 year, 10 months ago

Selected Answer: A

A is the right answer

upvoted 3 times

Bmaster 1 year, 10 months ago

Correct A.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #579

Topic 1


A company runs an application that uses Amazon RDS for PostgreSQL. The application receives traffic only on weekdays during business hours. The company wants to optimize costs and reduce operational overhead based on this usage.

Which solution will meet these requirements?

- A. Use the Instance Scheduler on AWS to configure start and stop schedules. **Most Voted**
- B. Turn off automatic backups. Create weekly manual snapshots of the database.
- C. Create a custom AWS Lambda function to start and stop the database based on minimum CPU utilization.
- D. Purchase All Upfront reserved DB instances.

Correct Answer: A

Community vote distribution



Comments

potomac **Highly Voted** 1 year, 1 month ago

Selected Answer: A

The Instance Scheduler on AWS solution automates the starting and stopping of Amazon Elastic Compute Cloud (Amazon EC2) and Amazon Relational Database Service (Amazon RDS) instances.

This solution helps reduce operational costs by stopping resources that are not in use and starting them when they are needed. The cost savings can be significant if you leave all of your instances running at full utilization continuously.

<https://aws.amazon.com/solutions/implementations/instance-scheduler-on-aws/>

upvoted 6 times

pentium75 **Most Recent** 11 months, 1 week ago

Selected Answer: A

B increases operational overhead

C Lambda functions could work but NOT "based on minimum CPU utilization"

D might save cost but not as much as A

upvoted 4 times

ibu007 1 year, 3 months ago

Selected Answer: A

A. Use the Instance Scheduler on AWS to configure start and stop schedules
upvoted 4 times

baba365 1 year, 2 months ago

Why not D?
upvoted 3 times

AndreiWebNet 1 year ago

How do you actually reduce costs enough to buy upfront instances that you pay for them if you use them 1 h or 24 and it is paid to run 24h. It says that you use this instances 8 hours a day 5 days a week, totaling 40h a week.... so is it the difference from 40h to 168 h?
upvoted 3 times

master9 11 months, 3 weeks ago

When you buy Reserved Instances, the larger the upfront payment, the greater the discount. To maximize your savings, you can pay all up-front and receive the largest discount. Partial up-front RI's offer lower discounts but give you the option to spend less up front. Lastly, you can choose to spend nothing up front and receive a smaller discount, but allowing you to free up capital to spend in other projects.

But you need some mechanism to stop on weekend and in night to save cost.
upvoted 2 times

ErnShm 1 year, 3 months ago

A
<https://docs.aws.amazon.com/solutions/latest/instance-scheduler-on-aws/solution-overview.html>
upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

Purpose-built scheduling minimizes operational overhead.
Aligns instance running time precisely with business hour demands.
Maintains backups unlike disabling auto backups.
More cost effective and flexible than reserved instances.
Simpler to implement than a custom Lambda function.
upvoted 4 times

anikety123 1 year, 3 months ago

Selected Answer: B

Its B. Check the AWS link

https://aws.amazon.com/solutions/implementations/instance-scheduler-on-aws/?nc1=h_ls
upvoted 1 times

anikety123 1 year, 3 months ago

Sorry I wanted to select A.
upvoted 5 times

mrsoa 1 year, 4 months ago

Selected Answer: A

A

<https://aws.amazon.com/solutions/implementations/instance-scheduler-on-aws/>
upvoted 2 times

luiscc 1 year, 4 months ago

Selected Answer: A

Scheduler do the job

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #580

Topic 1

A company uses locally attached storage to run a latency-sensitive application on premises. The company is using a lift and shift method to move the application to the AWS Cloud. The company does not want to change the application architecture.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure an Auto Scaling group with an Amazon EC2 instance. Use an Amazon FSx for Lustre file system to run the application.
- B. Host the application on an Amazon EC2 instance. Use an Amazon Elastic Block Store (Amazon EBS) GP2 volume to run the application.
- C. Configure an Auto Scaling group with an Amazon EC2 instance. Use an Amazon FSx for OpenZFS file system to run the application.
- D. Host the application on an Amazon EC2 instance. Use an Amazon Elastic Block Store (Amazon EBS) GP3 volume to run the application. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

TariqKipkemei **Highly Voted** 6 months, 2 weeks ago

Selected Answer: D

MOST cost-effectively =GP3
upvoted 5 times

potomac **Most Recent** 7 months ago

Selected Answer: D

gp3 offers SSD-performance at a 20% lower cost per GB than gp2 volumes.
upvoted 3 times

bojila 9 months, 1 week ago

GP3 is the latest version

upvoted 3 times

Hades2231 9 months, 1 week ago

Selected Answer: D

GP3 is the latest version, and it is cost effective

upvoted 3 times

Guru4Cloud 9 months, 3 weeks ago

Selected Answer: D

GP3 is preferable over GP2, FSx for Lustre, and FSx for OpenZFS is clear and convincing:

GP3 offers identical latency performance to GP2 at a lower price point.

FSx options are higher performance but more expensive and require application changes.

GP3 aligns better with lift and shift needs as a directly attached block storage volume.

upvoted 4 times

taustin2 9 months, 4 weeks ago

Selected Answer: D

Migrate your Amazon EBS volumes from gp2 to gp3 and save up to 20% on costs.

upvoted 3 times

Vadbro7 9 months, 4 weeks ago

Y not gp2

upvoted 1 times

Ale1973 10 months ago

Selected Answer: D

My rational:

Options A y C are based on autoscaling-group and no make sense for me on this scenary.

Then, use Amazon EBS is the solution and GP2 or GP3 is the question.

Requirement requires the most COST effective solution, then, I choose GP3

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #581

Topic 1

A company runs a stateful production application on Amazon EC2 instances. The application requires at least two EC2 instances to always be running.


A solutions architect needs to design a highly available and fault-tolerant architecture for the application. The solutions architect creates an Auto Scaling group of EC2 instances.

Which set of additional steps should the solutions architect take to meet these requirements?

- A. Set the Auto Scaling group's minimum capacity to two. Deploy one On-Demand Instance in one Availability Zone and one On-Demand Instance in a second Availability Zone.
- B. Set the Auto Scaling group's minimum capacity to four. Deploy two On-Demand Instances in one Availability Zone and two On-Demand Instances in a second Availability Zone. **Most Voted**
- C. Set the Auto Scaling group's minimum capacity to two. Deploy four Spot Instances in one Availability Zone.
- D. Set the Auto Scaling group's minimum capacity to four. Deploy two On-Demand Instances in one Availability Zone and two Spot Instances in a second Availability Zone.

Correct Answer: B

Community vote distribution



Comments

luiscc **Highly Voted** 1 year, 10 months ago

Selected Answer: B

By setting the Auto Scaling group's minimum capacity to four, the architect ensures that there are always at least two running instances. Deploying two On-Demand Instances in each of two Availability Zones ensures that the application is highly available and fault-tolerant. If one Availability Zone becomes unavailable, the application can still run in the other Availability Zone.

upvoted 19 times

Ale1973 **Highly Voted** 1 year, 10 months ago

Selected Answer: A

My rational is: Highly available = 2 AZ. and then 2 EC2 instances always running is 1 EC2 in each AZ. If an entire AZ fails.

my recommendation is to have 2 EC2 instances always running in each AZ. In an extreme case, SacalinGroup deploy the minimum instances (2) on the running AZ

upvoted 14 times

baba365 1 year, 8 months ago

Ans: A.

The application requires at least two EC2 instances to always be running = 2 minimum capacity... minimum cap of 4 ec2 will work but a waste of resources that doesn't follow well archi. framework.

upvoted 2 times

Ramdi1 1 year, 8 months ago

it says always have to have two running, hence you need 4. two in each AV. it might be a waste of resource but if that what is required by the company then so be it. Also you out the 4 you cannot use spot instances because if the two instances on the on demand go down and you need to use the spot instance they could be called back at any point.

upvoted 5 times

Ramdi1 1 year, 8 months ago

AZ * not AV

upvoted 3 times

bora4motion Most Recent 1 month, 1 week ago

Selected Answer: B

A and C are not compatible with the "fault tolerant" condition.

upvoted 1 times

emakid 11 months, 2 weeks ago

Option A:

Set the Auto Scaling group's minimum capacity to two. Deploy one On-Demand Instance in one Availability Zone and one On-Demand Instance in a second Availability Zone.

This configuration ensures that you have two instances running across two different AZs, which provides high availability. However, it does not take advantage of additional capacity to handle failures or spikes in demand. If either AZ becomes unavailable, you will have one running instance, but this does not meet the requirement of having at least two running instances at all times.

upvoted 1 times

emakid 11 months, 2 weeks ago

Selected Answer: B

Set the Auto Scaling group's minimum capacity to four. Deploy two On-Demand Instances in one Availability Zone and two On-Demand Instances in a second Availability Zone.

This configuration provides high availability with four instances distributed across two AZs. The minimum capacity of four ensures that even if one instance fails, there are still two instances in each AZ to handle the load. This option is highly available and fault-tolerant but may be more than required if only two instances are needed to be always running.

upvoted 2 times

cjace 12 months ago

Answer is A: If one Availability Zone fails, the Auto Scaling group will automatically launch a new instance in a different, healthy Availability Zone to maintain the desired capacity of two instances. This is one of the key benefits of using Auto Scaling groups—they automatically maintain the desired number of instances across multiple Availability Zones, ensuring that your application is highly available and fault-tolerant. So even in the event of a failure in one Availability Zone, your application will continue to run on the required number of instances. This is why it's recommended to distribute instances across multiple Availability Zones when designing architectures for high availability and fault tolerance.

upvoted 1 times

LuongTo 7 months ago

that is reasonable. but what do you think about "always"? With A, there's a moment there would be only one EC2 healthy.

upvoted 2 times

Marco_St 1 year, 4 months ago

Selected Answer: B

Selected Answer: B

so indeed ASG can set up a new EC2 instance in another AZ if there is one AZ failed with fault but it failed to meet the need of always having 2 instance running before the new instance replacement is done in the working AZ. so this is why we deploy 2 instances per AZ

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: B

If it would not mention the "stateful" application, and if it would only have to be "highly available" but NOT "fault-tolerant", A would be fine.

upvoted 5 times

1rob 1 year, 6 months ago

Selected Answer: B

From <<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/spot-best-practices.html> : Spot Instances are not suitable for workloads that are inflexible, stateful, fault-intolerant, or tightly coupled between instance nodes. So C and D don't fit.

From <<https://docs.aws.amazon.com/whitepapers/latest/real-time-communication-on-aws/use-multiple-availability-zones.html> : Within the constructs of AWS, customers are encouraged to run their workloads in more than one Availability Zone. This ensures that customer applications can withstand even a complete Availability Zone failure - a very rare event in itself.

So a HA solution in this case implies a total of 4 instances, 2 per AZ.

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

The main requirement here is a 'highly available and fault-tolerant architecture for the application', this covered by option B. The application requires at least two EC2 instances to always be running, main word here being 'atleast' which means more than two is ok.

upvoted 2 times

Ramdi1 1 year, 8 months ago

Selected Answer: B

B - Need 2 in each AZ and you cant use spot instances as it could be recalled.

upvoted 2 times

Mandar15 1 year, 8 months ago

Selected Answer: B

Stateful is keyword here. 2 is minimum required all time.

upvoted 2 times

JA2018 6 months, 2 weeks ago

minimum of 2x EC2 instances per AZ at all times.

upvoted 1 times

MII1975 1 year, 9 months ago

Selected Answer: A

If a complete AZ fails, autoscale will lunch a second EC2 in the running AZ. If that short period of time is not always, which is not, then the answer is B, but I would take my chances and select A in the exam xD because the application is highly available and fault-tolerant.

upvoted 1 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

° Minimum of 4 ensures at least 2 instances are always running in each AZ, meeting the HA requirement.

° On-Demand instances provide consistent performance and availability, unlike Spot.

° Spreading across 2 AZs adds fault tolerance, protecting from AZ failure.

upvoted 4 times

darkknight23 1 year, 10 months ago

burning123 1 year, 10 months ago

Selected Answer: B

While Spot Instances can be used to reduce costs, they might not provide the same level of availability and guaranteed uptime that On-Demand Instances offer. So I will go with B and not D.

upvoted 2 times

Sat897 1 year, 10 months ago

Selected Answer: B

Highly available - 2 AZ and then 2 EC2 instances always running. 2 in each AZ.

upvoted 3 times

Sat897 1 year, 10 months ago

Highly available - 2 AZ and then 2 EC2 instances always running. 2 in each AZ..

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #582

Topic 1


An ecommerce company uses Amazon Route 53 as its DNS provider. The company hosts its website on premises and in the AWS Cloud. The company's on-premises data center is near the us-west-1 Region. The company uses the eu-central-1 Region to host the website. The company wants to minimize load time for the website as much as possible.

Which solution will meet these requirements?

- A. Set up a geolocation routing policy. Send the traffic that is near us-west-1 to the on-premises data center. Send the traffic that is near eu-central-1 to eu-central-1. **Most Voted**
- B. Set up a simple routing policy that routes all traffic that is near eu-central-1 to eu-central-1 and routes all traffic that is near the on-premises datacenter to the on-premises data center.
- C. Set up a latency routing policy. Associate the policy with us-west-1.
- D. Set up a weighted routing policy. Split the traffic evenly between eu-central-1 and the on-premises data center.

Correct Answer: A

Community vote distribution



Comments

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: A

The key reasons are:

Geolocation routing allows you to route users to the closest endpoint based on their geographic location. This will provide the lowest latency.

Routing us-west-1 traffic to the on-premises data center minimizes latency for those users since it is also located near there.

Routing eu-central-1 traffic to the eu-central-1 AWS region minimizes latency for users nearby.

This achieves routing users to the closest endpoint on a geographic basis to optimize for low latency.

upvoted 5 times

PLN6302 1 year, 3 months ago

why can't be the option C

upvoted 2 times

lemur88 1 year, 3 months ago

You cannot associate the policy to us-west-1 as the AWS account is in eu-central-1
upvoted 5 times

awsgeek75 Most Recent 11 months ago

Selected Answer: A

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geo.html>

B can be done but definition of "near" is ambiguous

C wrong region

D wrong solution as splitting evenly does not reduce latency for on-prem server users

upvoted 2 times

Cyberkayu 11 months, 3 weeks ago

Selected Answer: A

not C. Client do not have AWS us-west-1 region. Client have a on prem DC near west-1

not D. 2 people visit the site together near eu-central-1, one of the user may be thrown to west-1 due to load balancing on split even weighted policy.

A and B are both valid, latency = how soon user reach the datacenter and received a responses from the DC, round trip. So in short, geolocation or send user to the nearest DC will improve latency.

upvoted 2 times

TariqKipkemei 1 year ago

Selected Answer: A

Geolocation routing policy allows you to route traffic based on the location of your users.

upvoted 4 times

t0nx 1 year ago

Selected Answer: C

C. Set up a latency routing policy. Associate the policy with us-west-1.

Explanation:

A latency routing policy directs traffic based on the lowest network latency to the specified AWS endpoint. Since the on-premises data center is near the us-west-1 Region, associating the policy with us-west-1 ensures that users near that region will be directed to the on-premises data center.

This allows for optimal routing, minimizing the load time for users based on their geographical proximity to the respective hosting locations (us-west-1 and eu-central-1).

Options A, B, and D do not explicitly consider latency or are not optimal for minimizing load time:

Option A (geolocation routing policy) would direct traffic based on the geographic location of the user but may not necessarily optimize for the lowest latency.

upvoted 2 times

awsgeek75 10 months, 3 weeks ago

There is nothing in us-west-1 as the company's data centre is near us-west-1.

upvoted 1 times

Chiquitabandita 1 year ago

except I don't think that it should be applied to the west region. If Geolocation is applied and the west is closer to the client, but the west is having intermittent issues at the time, they will have a longer latency even though closer to that region. this is why I would apply latency in a real world solution.

upvoted 1 times

Chiquitabandita 1 year ago

in real world I think it should use latency routing if the main concern is to lower the latency but AWS likes to promote geolocation and if that is in the question I think that will be the answer so I choose A.

upvoted 1 times

LeLe325 1 year, 9 months ago

Dada365 1 year, 2 months ago

The company wants to minimize load time for the website as much as possible... between data Centre and website or between users and website?

upvoted 1 times

Hades2231 1 year, 3 months ago

Selected Answer: A

Geolocation is the key word

upvoted 2 times

lemur88 1 year, 3 months ago

Selected Answer: A

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geo.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #583

Topic 1

A company has 5 PB of archived data on physical tapes. The company needs to preserve the data on the tapes for another 10 years for compliance purposes. The company wants to migrate to AWS in the next 6 months. The data center that stores the tapes has a 1 Gbps uplink internet connectivity.

Which solution will meet these requirements MOST cost-effectively?

- A. Read the data from the tapes on premises. Stage the data in a local NFS storage. Use AWS DataSync to migrate the data to Amazon S3 Glacier Flexible Retrieval.
- B. Use an on-premises backup application to read the data from the tapes and to write directly to Amazon S3 Glacier Deep Archive.
- C. Order multiple AWS Snowball devices that have Tape Gateway. Copy the physical tapes to virtual tapes in Snowball. Ship the Snowball devices to AWS. Create a lifecycle policy to move the tapes to Amazon S3 Glacier Deep Archive. **Most Voted**
- D. Configure an on-premises Tape Gateway. Create virtual tapes in the AWS Cloud. Use backup software to copy the physical tape to the virtual tape.

Correct Answer: C

Community vote distribution



Comments

adeyinkaamole **Highly Voted** 1 year, 9 months ago

If you have made it to the end of the exam dump, you will definitely pass your exams in Jesus name. After over a year of Procrastination, I am finally ready to write my AWS Solutions Architect Exam. Thank you Exam Topics
upvoted 31 times

Hades2231 **Highly Voted** 1 year, 9 months ago

Selected Answer: C

Ready for the exam tomorrow. Wish you guys all the best. BTW Snowball Device comes in handy when you need to move a huge amount of data but cant afford any bandwidth loss
upvoted 12 times

MatAlves **Most Recent** 8 months, 3 weeks ago

Oh, to think now we have to study 904 questions instead of just 583 lol
upvoted 3 times

JA2018 6 months, 2 weeks ago

currently (22 Nov 2024), question count now stands at 1.019
upvoted 3 times

XXXXXINN 8 months, 1 week ago

well, 2 weeks later today, it says 981 questions...
upvoted 3 times

awsgeek75 1 year, 4 months ago

Selected Answer: C

5PB over 1GB connection will take approximately 15 months so anything with "transfer" is invalid. ABD are not practical.
C: Just order snowball
upvoted 5 times

pentium75 1 year, 5 months ago

Selected Answer: C

Though we'll need more than 60 Snowball devices, C is the only option that works. The internet uplink could transport less than 2 PB in 6 months (otherwise, say with a 10 Gb uplink, D would work).
upvoted 6 times

Cyberkayu 1 year, 5 months ago

transfer 5 PB data in 1Gbps link, assume 0 overhead and drop packet, need 485 days, 10 hours, 50 minutes, 40 seconds to complete.

Snowball it is. C
upvoted 1 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

C
<https://docs.aws.amazon.com/storagegateway/latest/tgw/using-tape-gateway-snowball.html>
upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

Migrate petabyte-scale data stored on physical tapes to AWS using AWS Snowball
<https://aws.amazon.com/snowball/#:~:text=Migrate-,petabyte%2Dscale,-data%20stored%20on>
upvoted 2 times

hungta 1 year, 6 months ago

Selected Answer: C

5 PB data is too huge for using 1Gbps uplink. With this uplink, it takes more than 1 year to migrate this data.
upvoted 2 times

baba365 1 year, 8 months ago

Answer: D for most cost effective.

If you are looking for a cost-effective, durable, long-term, offsite alternative for data archiving, deploy a Tape Gateway. With its virtual tape library (VTL) interface, you can use your existing tape-based backup software infrastructure to store data on virtual tape cartridges that you create -

<https://docs.aws.amazon.com/storagegateway/latest/tgw/WhatIsStorageGateway.html>
upvoted 1 times

Devsin2000 1 year, 8 months ago

D
<https://aws.amazon.com/storagegateway/vtl/>

the bandwidth and available time is ample
upvoted 1 times

nnecode 1 year, 8 months ago

Selected Answer: A

The most cost-effective solution to meet the requirements is to read the data from the tapes on premises. Stage the data in a local NFS storage. Use AWS DataSync to migrate the data to Amazon S3 Glacier Flexible Retrieval.

This solution is the most cost-effective because it uses the least amount of bandwidth. AWS DataSync is a service that transfers data between on-premises storage and Amazon S3. It uses a variety of techniques to optimize the transfer speed and reduce c

upvoted 1 times

lemur88 1 year, 9 months ago

Selected Answer: C

Only thing that makes sense given the 1Gbps limitation

upvoted 1 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

Option C is likely the most cost-effective solution given the large data size and limited internet bandwidth. The physical data transfer and integration with the existing tape infrastructure provides efficiency benefits that can optimize the cost.

upvoted 3 times

barracouto 1 year, 9 months ago

Selected Answer: C

Went through this dump twice now. Exam is in about an hour. Will update with results.

upvoted 2 times

Vaishali12 1 year, 9 months ago

how was ur exam?
was these dump que helpful?

upvoted 1 times

riccardoto 1 year, 10 months ago

Finished the dump today - taking my exam tomorrow :) Wish me luck!

upvoted 4 times

Ale1973 1 year, 10 months ago

My rational: question is about which solution will meet these requirements MOST cost-effectively, not MOST time or effectively, then, my response is D (using Tape Gateways)

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #584

Topic 1

A company is deploying an application that processes large quantities of data in parallel. The company plans to use Amazon EC2 instances for the workload. The network architecture must be configurable to prevent groups of nodes from sharing the same underlying hardware.

Which networking solution meets these requirements?

- A. Run the EC2 instances in a spread placement group. **Most Voted**
- B. Group the EC2 instances in separate accounts.
- C. Configure the EC2 instances with dedicated tenancy.
- D. Configure the EC2 instances with shared tenancy.

Correct Answer: A*Community vote distribution***Comments****czyboi** **Highly Voted** 1 year, 9 months ago**Selected Answer: A**

A spread placement group is a group of instances that are each placed on distinct hardware.
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>

upvoted 12 times

pentium75 **Highly Voted** 1 year, 5 months ago**Selected Answer: A**

A, spread placement group does exactly what is required here.
Not C and D, tenancy determines whether the hardware is shared with other customers or not, it has nothing to do with your own instances sharing hardware. (On the contrary, dedicated tenancy would spread your EC2 instances across as little nodes as possible.)

Not B, accounts have nothing to do with the issue.

upvoted 5 times

emakid **Most Recent** 11 months, 2 weeks ago**Selected Answer: A**

Option A: Run the EC2 instances in a spread placement group.

Spread Placement Group: This placement group strategy ensures that EC2 instances are distributed across distinct hardware to reduce the risk of correlated failures. Instances in a spread placement group are placed on different underlying hardware, which aligns with the requirement to prevent groups of nodes from sharing the same underlying hardware. This is a good fit for the scenario where you need to ensure high availability and fault tolerance.

upvoted 2 times

Gape4 11 months, 2 weeks ago

Selected Answer: A

Spread – Strictly places a small group of instances across distinct underlying hardware to reduce correlated failures.

upvoted 2 times

Marco_St 1 year, 4 months ago

Selected Answer: A

dedicated tenancy cannot ensure the instances share the same hardware. So A

upvoted 3 times

maged123 1 year, 5 months ago

Selected Answer: A

Let's assume that you have two groups of instances, group A and group B and you have two physical hardware X and Y. With spread placement group, you can have group A of instances on hardware X and group B on hardware Y but this will not prevent hardware X to host other instances of other customers because your only requirement is to separate group A from group B. On the other hand, the dedicated tenancy means that AWS will dedicate the physical hardware only for you. So, the correct answer is A.

upvoted 3 times

Murtadhaceit 1 year, 5 months ago

Question is ambiguous and confusing. Is it asking about the EC2 instance of the same application not sharing hardware? or EC2 instance not sharing hardware with other EC2 from other applications?

upvoted 2 times

Mikado211 1 year, 6 months ago

Selected Answer: A

Spread placement group allows you to isolate your instances on hardware level.
Dedicated tenancy allows you to be sure that you are the only customer on the hardware.

The correct answer is A.

upvoted 2 times

Mikado211 1 year, 6 months ago

A : Spread placement group

upvoted 2 times

lucasbg 1 year, 6 months ago

Selected Answer: A

Def is A: <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

Keywords 'prevent groups of nodes from sharing the same underlying hardware'.

Spread Placement Group strictly places a small group of instances across distinct underlying hardware to reduce correlated failures.

upvoted 2 times

cciesam 1 year, 6 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>

Each instances is placed on seven different racks, each rack has its own network and power source.

upvoted 2 times

wsdasdasdqwdaw 1 year, 7 months ago

Another tricky question, but I would go for A because:

Dedicated instances:

Dedicated Instances are EC2 instances that run on hardware that's dedicated to a single customer. Dedicated Instances that belong to different AWS accounts are physically isolated at a hardware level, even if those accounts are linked to a single payer account. However, Dedicated Instances might share hardware with other instances from the same AWS account that are not Dedicated Instances.

Which is not the desired option.

Spread – strictly places a small group of instances across distinct underlying hardware to reduce correlated failures.

That's why A.

upvoted 3 times

garuta 1 year, 8 months ago

Selected Answer: C

C is clear.

upvoted 1 times

pentium75 1 year, 5 months ago

"Dedicated tenancy" means that all your nodes run on hardware that is not shared with other customers. This is counter-productive to the objective here.

upvoted 2 times

Devsin2000 1 year, 8 months ago

A

When you launch a new EC2 instance, the EC2 service attempts to place the instance in such a way that all of your instances are spread out across underlying hardware to minimize correlated failures.

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>

upvoted 3 times

taustin2 1 year, 8 months ago

Selected Answer: A

Spread Placement Group strictly places a small group of instances across distinct underlying hardware to reduce correlated failures.

upvoted 1 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: C

C is the correct answer.

Configuring the EC2 instances with dedicated tenancy ensures that each instance will run on isolated, single-tenant hardware. This meets the requirement to prevent groups of nodes from sharing underlying hardware.

A spread placement group only provides isolation at the Availability Zone level. Instances could still share hardware within an AZ.

upvoted 5 times

pentium75 1 year, 5 months ago

No. C ensures that your EC2 instances run on hardware that is not shared with other customers (!). It is still shared among YOUR instances.

upvoted 6 times

Eminenza22 1 year, 9 months ago

Selected Answer: A

Option A is the correct answer. It suggests running the EC2 instances in a spread placement group. This solution is cost-effective

Option A is the correct answer. It suggests running the EC2 instances in a spread placement group. This solution is cost effective and requires minimal development effort .

upvoted 3 times

Eminenza22 1 year, 9 months ago

The placement group reduces the risk of simultaneous failures by spreading the instances across distinct underlying hardware

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #585

Topic 1


A solutions architect is designing a disaster recovery (DR) strategy to provide Amazon EC2 capacity in a failover AWS Region. Business requirements state that the DR strategy must meet capacity in the failover Region.

Which solution will meet these requirements?

- A. Purchase On-Demand Instances in the failover Region.
- B. Purchase an EC2 Savings Plan in the failover Region.
- C. Purchase regional Reserved Instances in the failover Region.
- D. Purchase a Capacity Reservation in the failover Region. **Most Voted**

Correct Answer: D

Community vote distribution

**Comments**

EllenLiu 5 months, 2 weeks ago

Selected Answer: D

it is called on-demand capacity reservation actually . no discount. must specify the following attributes:

Availability Zone & Instance type & os type & tenancy

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html>

upvoted 2 times

EllenLiu 5 months, 2 weeks ago

DR is a classic use case for capacity reservation

upvoted 2 times

emakid 11 months, 2 weeks ago

Selected Answer: D

Option D: Purchase a Capacity Reservation in the failover Region.

Capacity Reservation: Capacity Reservations ensure that you have reserved capacity in a specific region for your instances, regardless of whether you are using On-Demand or Reserved Instances. This is ideal for DR scenarios because it guarantees that the required EC2 capacity will be available when needed.

upvoted 3 times

awsgeek75 1 year, 4 months ago

Selected Answer: D

"Business requirements state that the DR strategy must meet capacity in the failover Region"
so only D meets these requirements

- A. No reservation of capacity
- B. Saving plans don't guarantee capacity
- C. Can be possible but it's like an active instance so doesn't really make sense

upvoted 3 times

awsgeek75 1 year, 4 months ago

Correction on C (I mixed it up with tenancy!). Reserved instances are not really for capacity, it's for type of instance which gives good discount but that is not required here.

upvoted 2 times

Derek_G 1 year, 5 months ago

Selected Answer: D

Purchase a Capacity Reservation in the failover Region:

A Capacity Reservation allows you to reserve a specific amount of EC2 instance capacity in a given region without purchasing specific instances. This reserved capacity is dedicated to your account and can be utilized for launching instances when needed. Capacity Reservations offer flexibility, allowing you to launch different instance types and sizes within the reserved capacity.

Purchase regional Reserved Instances in the failover Region:

Regional Reserved Instances involve paying an upfront fee to reserve a certain number of specific EC2 instances in a particular region. These reserved instances are of a predefined type and size, providing a more traditional reservation model. Regional Reserved Instances are specific to a designated region and ensure that the reserved instances of a particular specification are available when needed.

upvoted 4 times

TheLaPlanta 1 year, 2 months ago

What I don't get is... can't you accomplish that by using on-demand? I understood that you can scale infinitely

upvoted 2 times

Tanidanindo 1 year, 2 months ago

It won't guarantee that you have the capacity when you need it. If available, like in most cases, it'll work. But the scenario requires a guarantee that the capacity will be available.

upvoted 2 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

D

Ask is to reserve capacity with RI capacity is not reserved also you can reserve capacity along with RI but only in AZ .

<https://repost.aws/knowledge-center/ri-reserved-capacity>

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: D

Capacity Reservations mitigate against the risk of being unable to get On-Demand capacity in case there are capacity constraints. If you have strict capacity requirements, and are running business-critical workloads that require a certain level of long or short-term capacity assurance, create a Capacity Reservation to ensure that you always have access to Amazon EC2 capacity when you need it, for as long as you need it.

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: D

Capacity Reservations enable you to reserve capacity for your Amazon EC2 instances in a specific Availability Zone for any duration. This gives you the flexibility to selectively add capacity reservations and still get the Regional RI discounts for that usage. By creating Capacity Reservations, you ensure that you always have access to Amazon EC2 capacity when you need it, for as long as you need it.

as long as you need it.

upvoted 3 times

potomac 1 year, 7 months ago

Savings Plans does not provide a capacity reservation.

upvoted 2 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: D

Capacity Reservations allocate EC2 capacity in a specific AWS Region for you to launch instances.

The capacity is reserved and available to be utilized when needed, meeting the requirement to provide EC2 capacity in the failover region.

Other options do not reserve capacity. On-Demand provides flexible capacity but does not reserve capacity upfront. Savings Plans and Reserved Instances provide discounts but do not reserve capacity.

Capacity Reservations allow defining instance attributes like instance type, platform, Availability Zone so the reserved capacity matches the production environment.

upvoted 4 times

Eminenza22 1 year, 9 months ago

Selected Answer: D

A regional Reserved Instance does not reserve capacity

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/reserved-instances-scope.html>

upvoted 3 times

judyda 1 year, 9 months ago

Selected Answer: D

reserved instances for price discount. need capacity reservation.

upvoted 3 times

gispankaj 1 year, 9 months ago

Selected Answer: C

The Reserved Instance discount applies to instance usage within the instance family, regardless of size.

upvoted 2 times

pentium75 1 year, 5 months ago

"Reserved Instances are not physical instances, but rather a billing discount "

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-reserved-instances.html>

upvoted 2 times

ErnShm 1 year, 9 months ago

D

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #586

Topic 1

A company has five organizational units (OUs) as part of its organization in AWS Organizations. Each OU correlates to the five businesses that the company owns. The company's research and development (R&D) business is separating from the company and will need its own organization. A solutions architect creates a separate new management account for this purpose.

What should the solutions architect do next in the new management account?

- A. Have the R&D AWS account be part of both organizations during the transition.
- B. Invite the R&D AWS account to be part of the new organization after the R&D AWS account has left the prior organization. **Most Voted**
- C. Create a new R&D AWS account in the new organization. Migrate resources from the prior R&D AWS account to the new R&D AWS account.
- D. Have the R&D AWS account join the new organization. Make the new management account a member of the prior organization.

Correct Answer: B

Community vote distribution

B (90%)

C (10%)

Comments

awsgeek75 **Highly Voted** 1 year, 4 months ago

Selected Answer: B

An account can only join another org when it leaves the first org.

A is wrong as it's not possible

C that's a new account so not really a migration

D The R&D department is separating from the company so you don't want the OU to join via nesting

upvoted 9 times

pentium75 **Highly Voted** 1 year, 5 months ago

Selected Answer: B

B as exactly described here: <https://repost.aws/knowledge-center/organizations-move-accounts>

upvoted 7 times

emakid **Most Recent** 11 months, 2 weeks ago

Selected Answer: B

Option B: Invite the R&D AWS account to be part of the new organization after the R&D AWS account has left the prior organization is the appropriate approach. This option ensures that the R&D AWS account transitions smoothly from the old organization to the new one. The steps involved are:

Remove the R&D AWS account from the existing organization: This is done from the existing organization's management account.

Invite the R&D AWS account to join the new organization: Once the R&D account is no longer part of the previous organization, it can be invited to and accepted into the new organization.

upvoted 2 times

Marco_St 1 year, 4 months ago

Selected Answer: B

<https://aws.amazon.com/blogs/mt/migrating-accounts-between-aws-organizations-with-consolidated-billing-to-all-features/>

upvoted 4 times

ale_brd_111 1 year, 5 months ago

Selected Answer: B

<https://repost.aws/knowledge-center/organizations-move-accounts>

Remove the member account from the old organization.

Send an invite to the member account from the new organization.

Accept the invite to the new organization from the member account.

upvoted 3 times

Derek_G 1 year, 5 months ago

Selected Answer: C

C is better. first migrate , then delete. avoid the data lose.

upvoted 1 times

pentium75 1 year, 5 months ago

What kind of "data lose" would happen when you change the account to a new organization? And why should you migrate ALL RESOURCES of the account to a new account?

upvoted 2 times

Derek_G 1 year, 5 months ago

C is better. first migrate , then delete. avoid the data lose.

upvoted 1 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

As per this document, B is clearly the answer.

<https://repost.aws/knowledge-center/organizations-move-accounts#:~:text=In%20either%20case%2C-perform%20these%20actions,-for%20each%20member>

upvoted 2 times

Joben 1 year, 8 months ago

Selected Answer: B

In either case, perform these actions for each member account:

- Remove the member account from the old organization.
- Send an invite to the member account from the new organization.
- Accept the invite to the new organization from the member account.

<https://repost.aws/knowledge-center/organizations-move-accounts>

upvoted 6 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: C

Creating a brand new AWS account in the new organization (Option C) allows for a clean separation and migration of only the

creating a brand new AWS account in the new organization (option C), ensuring a clean separation and migration of only the necessary resources from the old account to the new.

upvoted 2 times

pentium75 1 year, 5 months ago

"A clean separation" is already existing, they have their own account. "Migration of only the necessary resources from the old account to the new" is not asked for. They have an account in an existing organization, they need their own organization, thus move the existing account to a new organisation (B), done.

upvoted 2 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: C

When separating a business unit from an AWS Organizations structure, best practice is to:

Create a new AWS account dedicated for the business unit in the new organization

Migrate resources from the old account to the new account

Remove the old account from the original organization

This allows a clean break between the organizations and avoids any linking between them after separation.

upvoted 1 times

pentium75 1 year, 5 months ago

Says who?

upvoted 2 times

ErnShm 1 year, 9 months ago

B

<https://aws.amazon.com/blogs/mt/migrating-accounts-between-aws-organizations-with-consolidated-billing-to-all-features/>

upvoted 3 times

gispankaj 1 year, 9 months ago

Selected Answer: B

account can leave current organization and then join new organization.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #587

Topic 1

A company is designing a solution to capture customer activity in different web applications to process analytics and make predictions. Customer activity in the web applications is unpredictable and can increase suddenly. The company requires a solution that integrates with other web applications. The solution must include an authorization step for security purposes.

Which solution will meet these requirements?

- A. Configure a Gateway Load Balancer (GWLB) in front of an Amazon Elastic Container Service (Amazon ECS) container instance that stores the information that the company receives in an Amazon Elastic File System (Amazon EFS) file system. Authorization is resolved at the GWLB.
- B. Configure an Amazon API Gateway endpoint in front of an Amazon Kinesis data stream that stores the information that the company receives in an Amazon S3 bucket. Use an AWS Lambda function to resolve authorization.
- C. Configure an Amazon API Gateway endpoint in front of an Amazon Kinesis Data Firehose that stores the information that the company receives in an Amazon S3 bucket. Use an API Gateway Lambda authorizer to resolve authorization.
Most Voted
- D. Configure a Gateway Load Balancer (GWLB) in front of an Amazon Elastic Container Service (Amazon ECS) container instance that stores the information that the company receives on an Amazon Elastic File System (Amazon EFS) file system. Use an AWS Lambda function to resolve authorization.

Correct Answer: C*Community vote distribution***Comments****ralfj** **Highly Voted** 1 year, 9 months ago**Selected Answer: C**

<https://docs.aws.amazon.com/apigateway/latest/developerguide/apigateway-use-lambda-authorizer.html>
upvoted 9 times

FlyingHawk **Most Recent** 4 months, 2 weeks ago**Selected Answer: C**

<https://docs.aws.amazon.com/firehose/latest/dev/basic-deliver.html>

upvoted 1 times

emakid 11 months, 2 weeks ago

Selected Answer: C

Option C: Configure an Amazon API Gateway endpoint in front of an Amazon Kinesis Data Firehose that stores the information that the company receives in an Amazon S3 bucket. Use an API Gateway Lambda authorizer to resolve authorization.

This solution meets the requirements in the following ways:

Handles Unpredictable Traffic: Amazon Kinesis Data Firehose can handle variable amounts of streaming data and automatically scales to accommodate sudden increases in traffic.

Integration with Web Applications: Amazon API Gateway provides a RESTful API endpoint for integrating with web applications.

Authorization: An API Gateway Lambda authorizer provides the necessary authorization step to secure API access.

Data Storage: Amazon Kinesis Data Firehose can deliver data directly to an Amazon S3 bucket for storage, making it suitable for long-term analytics and predictions.

upvoted 4 times

Matte_ 1 year ago

Selected Answer: B

<https://docs.aws.amazon.com/apigateway/latest/developerguide/apigateway-use-lambda-authorizer.html>

upvoted 2 times

MatAlves 8 months, 3 weeks ago

You cannot use Kinesis Data stream to store data in S3. You need Firehose for that.

upvoted 3 times

Mr_Marcus 1 year ago

Ummm. This link (Use API Gateway Lambda authorizers) helps to validate "C" as the correct answer, not "B".

upvoted 2 times

4fad2f8 1 year, 4 months ago

Selected Answer: B

B. Amazon Kinesis Data Firehose does not save anything

upvoted 2 times

FlyingHawk 4 months, 2 weeks ago

Amazon Data Firehose can send data to S3: <https://docs.aws.amazon.com/firehose/latest/dev/s3-object-name.html>

upvoted 1 times

jaswantn 1 year, 4 months ago

option C...Amazon Kinesis Data Firehose that stores the information (that the company receives) in an Amazon S3 bucket. This answer statement is worded in a complex way. It means to say that Firehose stores the data in S3 ...which company receives from API Gateway.

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

Configure an Amazon API Gateway endpoint in front of an Amazon Kinesis Data Firehose that stores the information that the company receives in an Amazon S3 bucket. Use an API Gateway Lambda authorizer to resolve authorization.

upvoted 4 times

wsdasdasdqwda 1 year, 7 months ago

Using ECS just to store the information is an overkill. So B or C then, lambda authoriser is the key word => C

upvoted 4 times

Eminenza22 1 year, 9 months ago

Selected Answer: C

<https://docs.aws.amazon.com/lambda/latest/dg/services-kinesisfirehose.html>

upvoted 3 times

ErnShm 1 year, 9 months ago

C

authorizer is configured for the method. If it is, API Gateway calls the Lambda function. The Lambda function authenticates the caller by means such as the following: Calling out to an OAuth provider to get an OAuth access token

upvoted 3 times

gispankaj 1 year, 9 months ago

Selected Answer: C

lambda authoriser seems to be logical solution.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #588

Topic 1

An ecommerce company wants a disaster recovery solution for its Amazon RDS DB instances that run Microsoft SQL Server Enterprise Edition. The company's current recovery point objective (RPO) and recovery time objective (RTO) are 24 hours.

Which solution will meet these requirements MOST cost-effectively?

- A. Create a cross-Region read replica and promote the read replica to the primary instance.
- B. Use AWS Database Migration Service (AWS DMS) to create RDS cross-Region replication.
- C. Use cross-Region replication every 24 hours to copy native backups to an Amazon S3 bucket.
- D. Copy automatic snapshots to another Region every 24 hours. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

TiagueteVital **Highly Voted** 1 year, 9 months ago

Selected Answer: D

Snapshots are always a cost-efficiency way to have a DR plan.

upvoted 5 times

awsgeek75 **Highly Voted** 1 year, 4 months ago

Selected Answer: D

Cross region data transfer is billable so think of smallest amount of data to transfer every 24 hours

upvoted 5 times

MatAlves **Most Recent** 8 months, 3 weeks ago

A, B, C => cross-region \$\$

D => copy snapshots -> most cost-effectively.

upvoted 2 times

emakid 11 months, 2 weeks ago

Selected Answer: D

Option D: Copy automatic snapshots to another Region every 24 hours.

Explanation: This option involves copying RDS automatic snapshots to another Region. It is a straightforward way to ensure that snapshots are available in the event of a disaster. Since RDS snapshots are typically incremental and copied periodically, this solution matches the 24-hour RPO requirement effectively and is cost-effective compared to maintaining constant cross-Region replication.

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: D

Amazon RDS creates and saves automated backups of your DB instance or Multi-AZ DB cluster during the backup window of your DB instance. RDS creates a storage volume snapshot of your DB instance, backing up the entire DB instance and not just individual databases. RDS saves the automated backups of your DB instance according to the backup retention period that you specify. If necessary, you can recover your DB instance to any point in time during the backup retention period.

upvoted 3 times

wsdasdasdqwda 1 year, 7 months ago

most cost-effective way is just copying the snapshot (24h delta in the storage). => D

upvoted 3 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: D

Dddddddddd

upvoted 4 times

Eminenza22 1 year, 9 months ago

Selected Answer: D

This is the most cost-effective solution because it does not require any additional AWS services. Amazon RDS automatically creates snapshots of your DB instances every hour. You can copy these snapshots to another Region every 24 hours to meet your RPO and RTO requirements.

The other solutions are more expensive because they require additional AWS services. For example, AWS DMS is a more expensive service than AWS RDS.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #589

Topic 1


A company runs a web application on Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer that has sticky sessions enabled. The web server currently hosts the user session state. The company wants to ensure high availability and avoid user session state loss in the event of a web server outage.

Which solution will meet these requirements?

- A. Use an Amazon ElastiCache for Memcached instance to store the session data. Update the application to use ElastiCache for Memcached to store the session state.
- B. Use Amazon ElastiCache for Redis to store the session state. Update the application to use ElastiCache for Redis to store the session state. **Most Voted**
- C. Use an AWS Storage Gateway cached volume to store session data. Update the application to use AWS Storage Gateway cached volume to store the session state.
- D. Use Amazon RDS to store the session state. Update the application to use Amazon RDS to store the session state.

Correct Answer: B

Community vote distribution



Comments

Guru4Cloud **Highly Voted** 1 year, 8 months ago

Selected Answer: B

The key points are:

ElastiCache Redis provides in-memory caching that can deliver microsecond latency for session data.

Redis supports replication and multi-AZ which can provide high availability for the cache.

The application can be updated to store session data in ElastiCache Redis rather than locally on the web servers.

If a web server fails, the user can be routed via the load balancer to another web server which can retrieve their session data from the highly available ElastiCache Redis cluster.

upvoted 9 times

czyboi **Highly Voted** 1 year, 9 months ago

Selected Answer: B

high availability => use redis instead of Elastic memcache

high availability => use Redis instead of ElastiCache memcached
upvoted 5 times

emakid Most Recent 11 months, 2 weeks ago

Selected Answer: B

Option B: Use Amazon ElastiCache for Redis to store the session state. Update the application to use ElastiCache for Redis to store the session state.

Explanation: Amazon ElastiCache for Redis is suitable for session state storage because Redis provides both in-memory data storage and persistence options. Redis supports features like replication, persistence, and high availability (through Redis Sentinel or clusters). This ensures that session state is preserved and available even if individual web servers fail.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: B

As Memcached is not HA

upvoted 4 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

A

As cache needs to be distributed as ALB is used.

upvoted 1 times

potomac 1 year, 7 months ago

Selected Answer: B

B is correct

upvoted 3 times

franbarberan 1 year, 8 months ago

Selected Answer: D

Elastic cache is Only for RDS

upvoted 3 times

pentium75 1 year, 5 months ago

Since when?

upvoted 4 times

gispankaj 1 year, 9 months ago

Selected Answer: B

redis is correct since it provides high availability and data persistence

upvoted 5 times

Eminenza22 1 year, 9 months ago

Selected Answer: B

B is the correct answer. It suggests using Amazon ElastiCache for Redis to store the session state. Update the application to use ElastiCache for Redis to store the session state. This solution is cost-effective and requires minimal development effort.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #590

Topic 1

A company migrated a MySQL database from the company's on-premises data center to an Amazon RDS for MySQL DB instance. The company sized the RDS DB instance to meet the company's average daily workload. Once a month, the database performs slowly when the company runs queries for a report. The company wants to have the ability to run reports and maintain the performance of the daily workloads.

Which solution will meet these requirements?

- A. Create a read replica of the database. Direct the queries to the read replica. **Most Voted**
- B. Create a backup of the database. Restore the backup to another DB instance. Direct the queries to the new database.
- C. Export the data to Amazon S3. Use Amazon Athena to query the S3 bucket.
- D. Resize the DB instance to accommodate the additional workload.

Correct Answer: A

Community vote distribution

A (100%)

Comments

TariqKipkemei 6 months, 2 weeks ago

Selected Answer: A

queries for reports = read replica
upvoted 3 times

Guru4Cloud 8 months, 3 weeks ago

Selected Answer: A

Create a read replica of the database. Direct the queries to the read replica.
upvoted 4 times

Eminenza22 9 months, 1 week ago

Selected Answer: A

This is the most cost-effective solution because it does not require any additional AWS services. A read replica is a copy of a database that is synchronized with the primary database. You can direct the queries for the report to the read replica, which will not affect the performance of the daily workloads

upvoted 4 times

TiagueteVital 9 months, 1 week ago

Selected Answer: A

Clearly the right choice, with a read replica all the queries needed for a report are done in the replica, leaving the primary on best performance for write

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #591

Topic 1

A company runs a container application by using Amazon Elastic Kubernetes Service (Amazon EKS). The application includes microservices that manage customers and place orders. The company needs to route incoming requests to the appropriate microservices.

Which solution will meet this requirement MOST cost-effectively?

A. Use the AWS Load Balancer Controller to provision a Network Load Balancer.


B. Use the AWS Load Balancer Controller to provision an Application Load Balancer. **Most Voted**

C. Use an AWS Lambda function to connect the requests to Amazon EKS.

D. Use Amazon API Gateway to connect the requests to Amazon EKS.

Correct Answer: B

Community vote distribution



Comments

awsgeek75 **Highly Voted** 1 year, 4 months ago

Selected Answer: B

"The company needs to route incoming requests to the appropriate microservices"

In Kubernetes world, this would be called an Ingress Service so it will need B

<https://kubernetes-sigs.github.io/aws-load-balancer-controller/v2.6/>

<https://kubernetes.io/docs/concepts/services-networking/ingress/>

upvoted 7 times

pentium75 **Highly Voted** 1 year, 5 months ago

Selected Answer: B

Not D because

- even with an API gateway you'd need an ALB or ELB (so B+D would work, but D alone does not)

- you would use AWS API Gateway Controller (not "Amazon API Gateway") to create the API Gateway

upvoted 5 times

pentium75 1 year, 5 months ago

<https://aws.amazon.com/blog/containers/microservices-development-using-aws-controllers-for-kubernetes-eks-and>

<https://aws.amazon.com/blogs/containers/microservices-development-using-aws-controllers-for-kubernetes-ack-and-amazon-eks-blueprints/>

<https://aws.amazon.com/blogs/containers/integrate-amazon-api-gateway-with-amazon-eks/>
upvoted 5 times

emakid Most Recent 11 months, 2 weeks ago

Selected Answer: B

Option B: Use the AWS Load Balancer Controller to provision an Application Load Balancer (ALB).

Explanation: The AWS Load Balancer Controller can provision ALBs, which operate at the application layer (Layer 7). ALBs support advanced routing capabilities such as routing based on HTTP paths or hostnames. This makes ALBs well-suited for routing requests to different microservices based on URL paths or domains. This approach integrates well with Kubernetes and is a common pattern for microservices architectures.

upvoted 2 times

[Removed] 1 year, 5 months ago

Selected Answer: B

ALB is cost-effectively

upvoted 4 times

Mikado211 1 year, 6 months ago

Selected Answer: B

ALB is considered as expensive than API Gateway, particularly on higher load.

If you do not need any specific functionalities of API Gateway so you must choose ALB because it will be cheaper.

upvoted 4 times

Mikado211 1 year, 6 months ago

ALB is considered as LESS expensive

upvoted 2 times

riyasara 1 year, 6 months ago

Selected Answer: B

API Gateway has a pricing model that includes a cost per API call, and depending on the volume of requests, this could potentially be more expensive than using an Application Load Balancer.

upvoted 4 times

1rob 1 year, 6 months ago

Selected Answer: B

Routing requests to the appr. microserv. can easily be done with ALB and ingress. The ingress handles routing rules to the micro.serv. With answer D you wil still need ALB or NLB as can be seen in the pics of <https://aws.amazon.com/blogs/containers/integrate-amazon-api-gateway-with-amazon-eks/> or <https://aws.amazon.com/blogs/containers/microservices-development-using-aws-controllers-for-kubernetes-ack-and-amazon-eks-blueprints/> so that is not the most cost-effectively.

upvoted 3 times

ale_brd_111 1 year, 5 months ago

yeah, I was going with D than checked and seems that you are right to deploy API gateway you still a LB

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: D

Both ALB and API gateway can be used to route traffic to the microservices, but the question seeks the most 'cost effective' option.

You are charged for each hour or partial hour that an Application Load Balancer is running, and the number of Load Balancer Capacity Units (LCU) used per hour.

With Amazon API Gateway, you only pay when your APIs are in use.

I say API gateway is the best option for this case.

upvoted 2 times

pentium75 1 year, 5 months ago

But you still need an ALB or ELB

<https://aws.amazon.com/blogs/containers/microservices-development-using-aws-controllers-for-kubernetes-ack-and-amazon-eks-blueprints/>

<https://aws.amazon.com/blogs/containers/integrate-amazon-api-gateway-with-amazon-eks/>

upvoted 2 times

t0nx 1 year, 6 months ago

Selected Answer: B

AWS Load Balancer Controller: The AWS Load Balancer Controller is a Kubernetes controller that makes it easy to set up an Application Load Balancer (ALB) or Network Load Balancer (NLB) for your Amazon EKS clusters. It simplifies the process of managing load balancers for applications running on EKS.

Application Load Balancer (ALB): ALB is a Layer 7 load balancer that is capable of routing requests based on content, such as URL paths or hostnames. This makes it suitable for routing requests to different microservices based on specific criteria.

Cost-Effectiveness: ALB is typically more cost-effective than an NLB, and it provides additional features at the application layer, which may be useful for routing requests to microservices based on specific conditions.

Option D: Amazon API Gateway is designed for creating, publishing, and managing APIs. While it can integrate with Amazon EKS, it may be more feature-rich and complex than needed for simple routing to microservices within an EKS cluster.

upvoted 4 times

potomac 1 year, 7 months ago

Selected Answer: D

API Gateway provides an entry point to your microservices.

<https://aws.amazon.com/blogs/containers/integrate-amazon-api-gateway-with-amazon-eks/>

upvoted 3 times

ccmc 1 year, 7 months ago

B is correct, it is a required before exposing through api gateway

upvoted 2 times

thanhnv142 1 year, 7 months ago

B: is correct.

For EKS, use application load balancer to expose microservices

upvoted 4 times

KhasDenis 1 year, 8 months ago

Selected Answer: B

Routing to ms in k8s -> Ingresses -> Ingress Controller -> AWS Load Balancer Controller <https://kubernetes-sigs.github.io/aws-load-balancer-controller/v2.6/>

upvoted 4 times

RDM10 1 year, 8 months ago

Microservices--> API--> API GW

upvoted 3 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: D

D. Use Amazon API Gateway to connect the requests to Amazon EKS.

upvoted 3 times

MII1975 1 year, 9 months ago

Selected Answer: D

API Gateway is a fully managed service that makes it easy for you to create, publish, maintain, monitor, and secure APIs at any scale. API Gateway provides an entry point to your microservices.

upvoted 1 times

Eminenza22 1 year, 9 months ago

Selected Answer: D

<https://aws.amazon.com/blogs/containers/microservices-development-using-aws-controllers-for-kubernetes-ack-and-amazon-eks-blueprints/>

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #592

Topic 1

A company uses AWS and sells access to copyrighted images. The company's global customer base needs to be able to access these images quickly. The company must deny access to users from specific countries. The company wants to minimize costs as much as possible.

Which solution will meet these requirements?

- A. Use Amazon S3 to store the images. Turn on multi-factor authentication (MFA) and public bucket access. Provide customers with a link to the S3 bucket.
- B. Use Amazon S3 to store the images. Create an IAM user for each customer. Add the users to a group that has permission to access the S3 bucket.
- C. Use Amazon EC2 instances that are behind Application Load Balancers (ALBs) to store the images. Deploy the instances only in the countries the company services. Provide customers with links to the ALBs for their specific country's instances.
- D. Use Amazon S3 to store the images. Use Amazon CloudFront to distribute the images with geographic restrictions. Provide a signed URL for each customer to access the data in CloudFront. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

ralfj **Highly Voted** 9 months, 1 week ago

Selected Answer: D

Use Cloudfront and geographic restriction
upvoted 5 times

TariqKipkemei **Most Recent** 6 months, 2 weeks ago

Selected Answer: D

Store images = Amazon S3
global customer base needs to be able to access these images quickly = Amazon CloudFront
deny access to users from specific countries = Amazon CloudFront geographic restrictions, signed URLs
upvoted 4 times

Samuel 6 months, 4 weeks ago

Guru4Cloud 8 months, 4 weeks ago

Selected Answer: D

D. Use Amazon S3 to store the images. Use Amazon CloudFront to distribute the images with geographic restrictions. Provide a signed URL for each customer to access the data in CloudFront.

upvoted 4 times

Colz 8 months, 4 weeks ago

Correct answer is D

upvoted 1 times

hubbabubba 9 months, 1 week ago

Selected Answer: D

answer is D

upvoted 2 times

Eminenza22 9 months, 1 week ago

Selected Answer: D

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/georestrictions.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #593

Topic 1


A solutions architect is designing a highly available Amazon ElastiCache for Redis based solution. The solutions architect needs to ensure that failures do not result in performance degradation or loss of data locally and within an AWS Region. The solution needs to provide high availability at the node level and at the Region level.

Which solution will meet these requirements?

- A. Use Multi-AZ Redis replication groups with shards that contain multiple nodes. **Most Voted**
- B. Use Redis shards that contain multiple nodes with Redis append only files (AOF) turned on.
- C. Use a Multi-AZ Redis cluster with more than one read replica in the replication group.
- D. Use Redis shards that contain multiple nodes with Auto Scaling turned on.

Correct Answer: A

Community vote distribution



Comments

TariqKipkemei **Highly Voted** 1 year ago

Selected Answer: A

Multi-AZ is only supported on Redis clusters that have more than one node in each shard (node groups).

<https://docs.aws.amazon.com/AmazonElastiCache/latest/red-ug/AutoFailover.html#:~:text=node%20in%20each-,shard,-Topics>
upvoted 5 times

pentium75 **Highly Voted** 11 months, 1 week ago

Selected Answer: A

It seems like "Multi-AZ Redis replication group" (A) and "Multi-AZ Redis cluster" (C) are different wordings for the same configuration. However, "to minimize the impact of a node failure, we recommend that your implementation use multiple nodes in each shard" - and that is mentioned only in A.

upvoted 5 times

LocNV **Most Recent** 11 months, 2 weeks ago

Selected Answer: A

High availability at the node level = shard and Multi-AZ = region level

high availability at the node level = shard and multi A-Z = region level

upvoted 5 times

Cyberkayu 11 months, 3 weeks ago

did client ask for improved performance, unfortunately they didn't, so C is good to have but not part of the business requirement.

My answer A.

upvoted 3 times

SHAAHIBHUSHANAWS 1 year ago

A

Multi-AZ is only option. It is regional service so can use backup to replicate but can not use for failover.

upvoted 2 times

t0nx 1 year ago

Selected Answer: C

C. Use a Multi-AZ Redis cluster with more than one read replica in the replication group.

In summary, option C, using a Multi-AZ Redis cluster with more than one read replica, is designed to provide both node-level and AWS Region-level high availability, making it the most suitable choice for the given requirements.

upvoted 2 times

potomac 1 year, 1 month ago

Selected Answer: A

the replication structure is contained within a shard (called node group in the API/CLI) which is contained within a Redis cluster

A shard (in the API and CLI, a node group) is a hierarchical arrangement of nodes, each wrapped in a cluster. Shards support replication. Within a shard, one node functions as the read/write primary node. All the other nodes in a shard function as read-only replicas of the primary node.

upvoted 1 times

thanhnv142 1 year, 1 month ago

C is correct.

Not A because in replication mode, shard have multiple nodes by default.

B and D not correct because that not an option

upvoted 1 times

iwannabeawsgod 1 year, 1 month ago

Selected Answer: C

its c for me

upvoted 1 times

bsbs1234 1 year, 1 month ago

C:

Cluster mode will create multiple shards, when node level failure, request of shard that not impacted will not has any performance impact. If the issue at AZ level, spread traffic between multiple shards shall also reduce the performance degrade.

upvoted 1 times

loveaws 1 year, 2 months ago

C.

Option A is not ideal because it doesn't mention read replicas, and it's generally better to have read replicas for both performance and high availability.

Option B mentions Redis append-only files (AOF), but AOF alone doesn't provide high availability or fault tolerance.

Option D mentions Auto Scaling, but this doesn't directly address high availability at the Region level or data replication

upvoted 1 times

taustin2 1 year, 2 months ago

Multi-AZ is only supported on Redis clusters that have more than one node in each shard.

upvoted 1 times

taustin2 1 year, 2 months ago

Selected Answer: C

Selected Answer: A

<https://docs.aws.amazon.com/AmazonElastiCache/latest/red-ug/Replication.html>

upvoted 5 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: A

Multi-AZ replication groups provide automatic failover between AZs if there is an issue with the primary AZ. This provides high availability at the region level

upvoted 3 times

pentium75 11 months, 1 week ago

What about "the node level"?

upvoted 1 times

xyb 1 year, 2 months ago

Selected Answer: C

Enabling ElastiCache Multi-AZ with automatic failover on your Redis cluster (in the API and CLI, replication group) improves your fault tolerance. This is true particularly in cases where your cluster's read/write primary cluster becomes unreachable or fails for any reason. Multi-AZ with automatic failover is only supported on Redis clusters that support replication

upvoted 1 times

MII1975 1 year, 3 months ago

Selected Answer: A

I would go with A too

I would go with A, Using AOF can't protect you from all failure scenarios.

For example, if a node fails due to a hardware fault in an underlying physical server, ElastiCache will provision a new node on a different server. In this case, the AOF is not available and can't be used to recover the data.

upvoted 2 times

hubbabubba 1 year, 3 months ago

Selected Answer: A

Hate to say this, but I read the two docs linked below, and I still think the answer is A. Turning on AOF helps in data persistence after failure, but it does nothing for availability unless you use Multi-AZ replica groups.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #594

Topic 1

A company plans to migrate to AWS and use Amazon EC2 On-Demand Instances for its application. During the migration testing phase, a technical team observes that the application takes a long time to launch and load memory to become fully productive.

Which solution will reduce the launch time of the application during the next testing phase?

- A. Launch two or more EC2 On-Demand Instances. Turn on auto scaling features and make the EC2 On-Demand Instances available during the next testing phase.
- B. Launch EC2 Spot Instances to support the application and to scale the application so it is available during the next testing phase.
- C. Launch the EC2 On-Demand Instances with hibernation turned on. Configure EC2 Auto Scaling warm pools during the next testing phase. **Most Voted**
- D. Launch EC2 On-Demand Instances with Capacity Reservations. Start additional EC2 instances during the next testing phase.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Guru4Cloud** **Highly Voted** 1 year, 2 months ago**Selected Answer: C**

Using EC2 hibernation and Auto Scaling warm pools will help address this:

Hibernation saves the in-memory state of the EC2 instance to persistent storage and shuts the instance down. When the instance is started again, the in-memory state is restored, which launches much faster than launching a new instance.

Warm pools pre-initialize EC2 instances and keep them ready to fulfill requests, reducing launch time. The hibernated instances can be added to a warm pool.

When auto scaling scales out during the next testing phase, it will be able to launch instances from the warm pool rapidly since they are already initialized

upvoted 8 times

awsgeek75 **Most Recent** 11 months ago**Selected Answer: C**

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/ec2-auto-scaling-warm-pools.html>

upvoted 3 times

riyasara 1 year ago

Selected Answer: C

Amazon EC2 hibernation and warm pool

upvoted 2 times

TariqKipkemei 1 year ago

Selected Answer: C

If an instance or application takes a long time to bootstrap and build a memory footprint in order to become fully productive, you can use hibernation to pre-warm the instance. To pre-warm the instance, you:

Launch it with hibernation enabled.

Bring it to a desired state.

Hibernate it so that it's ready to be resumed to the desired state whenever needed.

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Hibernate.html#:~:text=you%20can%20use-,hibernation,-to%20pre%2Dwarm>

upvoted 4 times

potomac 1 year, 1 month ago

Selected Answer: C

With Amazon EC2 hibernation enabled, you can maintain your EC2 instances in a "pre-warmed" state so these can get to a productive state faster.

upvoted 3 times

tabbyDolly 1 year, 2 months ago

C: <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Hibernate.html>

upvoted 3 times

ralfj 1 year, 3 months ago

Selected Answer: C

just use hibernation option so you won't load the full EC2 Instance

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #595

Topic 1

A company's applications run on Amazon EC2 instances in Auto Scaling groups. The company notices that its applications experience sudden traffic increases on random days of the week. The company wants to maintain application performance during sudden traffic increases.

Which solution will meet these requirements MOST cost-effectively?

- A. Use manual scaling to change the size of the Auto Scaling group.
- B. Use predictive scaling to change the size of the Auto Scaling group.
- C. Use dynamic scaling to change the size of the Auto Scaling group. **Most Voted**
- D. Use schedule scaling to change the size of the Auto Scaling group.

Correct Answer: C

Community vote distribution



Comments

ralfj **Highly Voted** 1 year, 3 months ago

Selected Answer: C

Dynamic Scaling – This is yet another type of Auto Scaling in which the number of EC2 instances is changed automatically depending on the signals received. Dynamic Scaling is a good choice when there is a high volume of unpredictable traffic.

<https://www.developer.com/web-services/aws-auto-scaling-types-best-practices/#:~:text=Dynamic%20Scaling%20%E2%80%93%20This%20is%20yet,high%20volume%20of%20unpredictable%20traffic>

upvoted 6 times

tabbyDolly **Highly Voted** 1 year, 2 months ago

C - " sudden traffic increases on random days of the week" --> dynamic scaling

upvoted 5 times

awsgeek75 **Most Recent** 10 months, 3 weeks ago

Selected Answer: C

random = dynamic

A: Manual is never a solution
B: Predictive is not possible as it's random
D: Cannot schedule random
upvoted 4 times

TariqKipkemei 1 year ago

Selected Answer: C

Dynamic scaling
upvoted 2 times

dilaaziz 1 year, 1 month ago

Selected Answer: C

<https://aws.amazon.com/ec2/autoscaling/faqs/>
upvoted 3 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: C

C is the best answer here. Dynamic scaling is the most cost-effective way to automatically scale the Auto Scaling group to maintain performance during random traffic spikes.
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #596

Topic 1


An ecommerce application uses a PostgreSQL database that runs on an Amazon EC2 instance. During a monthly sales event, database usage increases and causes database connection issues for the application. The traffic is unpredictable for subsequent monthly sales events, which impacts the sales forecast. The company needs to maintain performance when there is an unpredictable increase in traffic.

Which solution resolves this issue in the MOST cost-effective way?

- A. Migrate the PostgreSQL database to Amazon Aurora Serverless v2. **Most Voted**
- B. Enable auto scaling for the PostgreSQL database on the EC2 instance to accommodate increased usage.
- C. Migrate the PostgreSQL database to Amazon RDS for PostgreSQL with a larger instance type.
- D. Migrate the PostgreSQL database to Amazon Redshift to accommodate increased usage.

Correct Answer: A

Community vote distribution



Comments

Guru4Cloud **Highly Voted** 1 year, 2 months ago

Selected Answer: A

Answer is A.

Aurora Serverless v2 got autoscaling, highly available and cheaper when compared to the other options.

upvoted 8 times

pentium75 **Most Recent** 11 months, 1 week ago

Selected Answer: A

Not B - we can auto-scale the EC2 instance, but not "the [self-managed] PostgreSQL database ON the EC2 instance"
 Not C - This does not mention scaling, so it would incur high cost and still it might not be able to keep up with the "unpredictable" spikes

Not D - Redshift is OLAP Data Warehouse

upvoted 3 times

TariqKipkemei 1 year ago

Selected Answer: A

Amazon Aurora Serverless is an on-demand, auto-scaling configuration for Aurora where the database automatically starts up, shuts down, and scales capacity up or down based on your application's needs. This is the least costly option for unpredictable traffic.

upvoted 3 times

tabbyDolly 1 year, 2 months ago

A: "he traffic is unpredictable for subsequent monthly sales events" --> serverless

upvoted 3 times

Wayne23Fang 1 year, 3 months ago

Selected Answer: C

A is probably more expensive than C. Aurora is serverless and fast. But nevertheless it needs DB migration service. Not sure DMS may not be free.

upvoted 1 times

danielmakita 1 year, 1 month ago

C is more expensive if you think the scenario where the traffic is low. You are paying for a larger hardware but not using it. That's why I think A is correct.

upvoted 4 times

TiagueteVital 1 year, 3 months ago

Selected Answer: A

A to autoscaling

upvoted 3 times

manOfThePeople 1 year, 3 months ago

Answer is A.

Aurora Serverless v2 got autoscaling, highly available and cheaper when compared to the other options.

upvoted 2 times

anikety123 1 year, 3 months ago

Selected Answer: A

The correct answer is A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #597

Topic 1

A company hosts an internal serverless application on AWS by using Amazon API Gateway and AWS Lambda. The company's employees report issues with high latency when they begin using the application each day. The company wants to reduce latency.

Which solution will meet these requirements?

- A. Increase the API Gateway throttling limit.
- B. Set up a scheduled scaling to increase Lambda provisioned concurrency before employees begin to use the application each day. **Most Voted**
- C. Create an Amazon CloudWatch alarm to initiate a Lambda function as a target for the alarm at the beginning of each day.
- D. Increase the Lambda function memory.

Correct Answer: B

Community vote distribution

B (100%)

Comments

oayoad **Highly Voted** 1 year, 9 months ago

Selected Answer: B

<https://aws.amazon.com/blogs/compute/scheduling-aws-lambda-provisioned-concurrency-for-recurring-peak-usage/>
upvoted 5 times

Guru4Cloud **Highly Voted** 1 year, 8 months ago

Selected Answer: B

Set up a scheduled scaling to increase Lambda provisioned concurrency before employees begin to use the application each day.

upvoted 5 times

emakid **Most Recent** 11 months, 2 weeks ago

Selected Answer: B

Option B: Set up a scheduled scaling to increase Lambda provisioned concurrency before employees begin to use the

Option B: Set up a scheduled scaling to increase Lambda provisioned concurrency before employees begin to use the application each day.

Explanation: Provisioned concurrency ensures that a specified number of Lambda instances are initialized and ready to handle requests. By scheduling this scaling, you can pre-warm Lambda functions before peak usage times, reducing cold start latency. This solution directly addresses the latency issue caused by cold starts.

upvoted 4 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

Provisioned concurrency pre-initializes execution environments for your functions. These execution environments are prepared to respond immediately to incoming function requests at start of day.

upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: B

A is wrong

API Gateway throttling limit is for better throughput, not for latency

upvoted 2 times

MI1975 1 year, 9 months ago

Selected Answer: B

Provisioned Concurrency incurs additional costs, so it is cost-efficient to use it only when necessary. For example, early in the morning when activity starts, or to handle recurring peak usage.

upvoted 4 times

Eminenza22 1 year, 9 months ago

Selected Answer: B

B option setting up a scheduled scaling to increase Lambda provisioned concurrency before employees begin to use the application each day. This solution is cost-effective and requires minimal development effort.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #598

Topic 1

A research company uses on-premises devices to generate data for analysis. The company wants to use the AWS Cloud to analyze the data. The devices generate .csv files and support writing the data to an SMB file share. Company analysts must be able to use SQL commands to query the data. The analysts will run queries periodically throughout the day.

Which combination of steps will meet these requirements MOST cost-effectively? (Choose three.)

A. Deploy an AWS Storage Gateway on premises in Amazon S3 File Gateway mode. **Most Voted**

B. Deploy an AWS Storage Gateway on premises in Amazon FSx File Gateway mode.

C. Set up an AWS Glue crawler to create a table based on the data that is in Amazon S3. **Most Voted**


D. Set up an Amazon EMR cluster with EMR File System (EMRFS) to query the data that is in Amazon S3. Provide access to analysts.

E. Set up an Amazon Redshift cluster to query the data that is in Amazon S3. Provide access to analysts.

F. Set up Amazon Athena to query the data that is in Amazon S3. Provide access to analysts. **Most Voted**

Correct Answer: ACF

Community vote distribution



Comments

awsgeek75 **Highly Voted** 1 year, 4 months ago

Selected Answer: ACF

SQL Queries is Athena so DE are wrong and we are now dependant on S3

A to get files into S3

C Glue to convert CSV to S3 table data

B irrelevant as we don't have anything to consume data from FSx in other options

upvoted 6 times

JA2018 6 months, 2 weeks ago

err,,, pardon my ignorance but I did not recall seeing any mention of S3 in the stem

upvoted 1 times

awsgeek75 1 year, 4 months ago

My only reservation with this answer is C.

CSV is technically a table and Athena can query multiple csv from S3. Glue just seems overengineering over here

upvoted 3 times

pentium75 Highly Voted 1 year, 5 months ago

Selected Answer: ACF

A to upload the files to S3 via SMB

C to convert the data from CSV format

F to query with SQL

Not B (we need the data in S3, not in FSx)

Not D or E (we should provide the ability to run SQL queries)

upvoted 5 times

TariqKipkemei Most Recent 1 year, 6 months ago

Selected Answer: ACF

SMB + use SQL commands to query the data = Amazon S3 File Gateway mode + Amazon Athena

upvoted 1 times

wsdasdasdqwda 1 year, 7 months ago

<https://aws.amazon.com/storagegateway/file/s3/#:~:text=Amazon%20S3%20File%20Gateway%20provides,Amazon%20S3%20with%20local%20caching.>

"Amazon S3 File Gateway provides a seamless way to connect to the cloud in order to store application data files and backup images as durable objects in Amazon S3 cloud storage. Amazon S3 File Gateway offers SMB or NFS-based access to data in Amazon S3 with local caching"

=> SMB and NFS is supported in Amazon S3 File Gateway => ACF

upvoted 2 times

iwannabeawsgod 1 year, 7 months ago

Selected Answer: ACF

ACF 100% sure

upvoted 3 times

Ramdi1 1 year, 8 months ago

Selected Answer: ACF

I thought the correct answer was BCF however I have changed my mind to BCF

FSx does support SMB protocol. However so does s3 file gateway which is version 2 and 3 of the SMB protocol. Hence using it with athena ACF should be correct

upvoted 4 times

RDM10 1 year, 8 months ago

SMB file share- is B incorrect?

upvoted 1 times

pentium75 1 year, 5 months ago

Yes because "FSx File Gateway" uploads the files to FSx, but we need it in S3.

upvoted 2 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: BCE

BCF is the correct

upvoted 1 times

pentium75 1 year, 5 months ago

No because FSx File Gateway (B) uploads it to FSx while we we need it in S3. S3 File Gateway provides access to S3 via SMB.

upvoted 3 times

Eminenza22 1 year, 9 months ago

Selected Answer: ACF

<https://docs.aws.amazon.com/glue/latest/dg/aws-glue-programming-etl-format-csv-home.html>
<https://aws.amazon.com/blogs/aws/amazon-athena-interactive-sql-queries-for-data-in-amazon-s3/>
<https://aws.amazon.com/storagegateway/faqs/>
upvoted 3 times

anikety123 1 year, 9 months ago

Selected Answer: ACF

It should be ACF
upvoted 2 times

ralfj 1 year, 9 months ago

Selected Answer: ACF

ACF use S3 File Gateway, Use Glue and Use Athena
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #599

Topic 1

A company wants to use Amazon Elastic Container Service (Amazon ECS) clusters and Amazon RDS DB instances to build and run a payment processing application. The company will run the application in its on-premises data center for compliance purposes.


A solutions architect wants to use AWS Outposts as part of the solution. The solutions architect is working with the company's operational team to build the application.

Which activities are the responsibility of the company's operational team? (Choose three.)

- A. Providing resilient power and network connectivity to the Outposts racks **Most Voted**
- B. Managing the virtualization hypervisor, storage systems, and the AWS services that run on Outposts
- C. Physical security and access controls of the data center environment **Most Voted**
- D. Availability of the Outposts infrastructure including the power supplies, servers, and networking equipment within the Outposts racks
- E. Physical maintenance of Outposts components
- F. Providing extra capacity for Amazon ECS clusters to mitigate server failures and maintenance events **Most Voted**

Correct Answer: ACF

Community vote distribution



Comments

taustin2 **Highly Voted** 1 year, 7 months ago

Selected Answer: ACF

From <https://docs.aws.amazon.com/whitepapers/latest/aws-outposts-high-availability-design/aws-outposts-high-availability-design.html>

With Outposts, you are responsible for providing resilient power and network connectivity to the Outpost racks to meet your availability requirements for workloads running on Outposts. You are responsible for the physical security and access controls of the data center environment. You must provide sufficient power, space, and cooling to keep the Outpost operational and network connections to connect the Outpost back to the Region. Since Outpost capacity is finite and determined by the size and number of racks AWS installs at your site, you must decide how much EC2, EBS, and S3 on Outposts capacity you need to

run your initial workloads, accommodate future growth, and to provide extra capacity to mitigate server failures and maintenance events.

upvoted 27 times

ibu007 Highly Voted 1 year, 9 months ago

Selected Answer: ACE

My exam is tomorrow. thank you all for the answers and links.

upvoted 16 times

pentium75 1 year, 5 months ago

"Physical maintenance" such as replacing faulty disks is NOT your responsibility.

upvoted 7 times

FlyingHawk Most Recent 4 months, 2 weeks ago

Selected Answer: ACF

this is a tricky question. I selected ACE initially, but after reading the AWS doc:

<https://docs.aws.amazon.com/whitepapers/latest/aws-outposts-high-availability-design/aws-outposts-high-availability-design.html>

Based on it , capacity is customer's responsibility. E is AWS responsibility

upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

Since Outpost capacity is finite and determined by the size and number of racks AWS installs at your site, you must decide how much EC2, EBS, and S3 on Outposts capacity you need to run your initial workloads, accommodate future growth, and to provide extra capacity to mitigate server failures and maintenance events.(F)

AWS is responsible for the availability of the Outposts infrastructure including the power supplies, servers, and networking equipment within the AWS Outposts racks. (D, E) AWS also manages the virtualization hypervisor, storage systems, and the AWS services that run on Outposts.(B)

upvoted 1 times

[Removed] 9 months, 4 weeks ago

Selected Answer: ACE

The activities that are the responsibility of the company's operational team when using Amazon Elastic Container Service (Amazon ECS) clusters and Amazon RDS DB instances on AWS Outposts are:

Providing resilient power and network connectivity to the Outposts racks.

Physical security and access controls of the data center environment.

Physical maintenance of Outposts components.

The solutions architect is responsible for the following activities:

Managing the virtualization hypervisor, storage systems, and the AWS services that run on Outposts.

Ensuring the availability of the Outposts infrastructure, including the power supplies, servers, and networking equipment within the Outposts racks.

Providing extra capacity for Amazon ECS clusters to mitigate server failures and maintenance events.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: ACF

F: "If there is no additional capacity on the Outpost, the instance remains in the stopped state. The Outpost owner can try to free up used capacity or request additional capacity for the Outpost so that the migration can complete."

Not D: "Equipment within the Outposts rack" is AWS' responsibility, you're not supposed to touch that

Not E: "When the AWS installation team arrives on site, they will replace the unhealthy hosts, switches, or rack elements"

upvoted 6 times

1rob 1 year, 6 months ago

Selected Answer: ACF

From <<https://aws.amazon.com/outposts/rack/faqs/>> : Your site must support the basic power, networking and space requirements to host an Outpost ==> A

From <<https://docs.aws.amazon.com/whitepapers/latest/applying-security-practices-to-network-workload-for-csps/the-shared-responsibility-model.html>> : In AWS Outposts. the customer takes the responsibility of securing the physical infrastructure

shared responsibility, management of the capacity, and customer takes the responsibility of securing the physical infrastructure to host the AWS Outposts equipment in their own data centers. ==> C

upvoted 2 times

1rob 1 year, 6 months ago

and From <<https://docs.aws.amazon.com/whitepapers/latest/aws-outposts-high-availability-design/aws-outposts-high-availability-design.html>> : Since Outpost capacity is finite and determined by the size and number of racks AWS installs at your site, you must decide how much EC2, EBS, and S3 on Outposts capacity you need to run your initial workloads, accommodate future growth, and to provide extra capacity to mitigate server failures and maintenance events. ==> F

upvoted 3 times

1rob 1 year, 6 months ago

From <<https://docs.aws.amazon.com/whitepapers/latest/aws-outposts-high-availability-design/aws-outposts-high-availability-design.html>>

AWS is responsible for the availability of the Outposts infrastructure including the power supplies, servers, and networking equipment within the AWS Outposts racks. AWS also manages the virtualization hypervisor, storage systems, and the AWS services that run on Outposts. So The customer isn't so not D.

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: AC

Only A and C are correct.

AWS is responsible for the hardware and software that run on AWS Outposts. This is a fully managed infrastructure service. AWS manages security patches, updates firmware, and maintains the Outpost equipment. AWS also monitors the performance, health, and metrics for your Outpost and determines whether any maintenance is required.

<https://docs.aws.amazon.com/outposts/latest/userguide/outpost-maintenance.html>

upvoted 1 times

pentium75 1 year, 5 months ago

"Choose three".

You're missing F, you must order the Outposts rack with excess capacity

upvoted 2 times

[Removed] 1 year, 6 months ago

Selected Answer: A,E

The role that physical companies will play is ACE.

upvoted 1 times

potomac 1 year, 7 months ago

Selected Answer: ACD

E is wrong

If there is a need to perform physical maintenance, AWS will reach out to schedule a time to visit your site.

<https://aws.amazon.com/outposts/rack/faqs/#:~:text=As%20AWS%20Outposts%20rack%20runs,the%20Outpost%20for%20compliance%20certification.>

upvoted 1 times

pentium75 1 year, 5 months ago

So is D, "equipment WITHIN the Outposts rack" is something that your Infra team should stay away from.

upvoted 2 times

beast2091 1 year, 7 months ago

ACE

AWS is responsible for the availability of the Outposts infrastructure including the power supplies, servers, and networking equipment within the AWS Outposts racks. AWS also manages the virtualization hypervisor, storage systems, and the AWS services that run on Outposts.

<https://d1.awsstatic.com/whitepapers/aws-outposts-high-availability-design-and-architecture-considerations.pdf>

upvoted 1 times

dilaaziz 1 year, 7 months ago

Selected Answer: ACF

<https://docs.aws.amazon.com/whitepapers/latest/aws-outposts-high-availability-design/aws-outposts-high-availability-design.html>

upvoted 2 times

canonlycontainletters1 1 year, 7 months ago

Selected Answer: ACD

I choose ACD

upvoted 1 times

danielmakita 1 year, 7 months ago

Selected Answer: ACD

I think ACD is correct

upvoted 1 times

chris0975 1 year, 7 months ago

Selected Answer: ACF

You get to choose the capacity. F

upvoted 2 times

thanhnv142 1 year, 7 months ago

A, C and D

upvoted 1 times

aleksand41 1 year, 8 months ago

ACD <https://docs.aws.amazon.com/outposts/latest/userguide/outpost-maintenance.html>

upvoted 1 times

Ramdi1 1 year, 8 months ago

Selected Answer: ACD

I think because of the shared responsibility model it is ACD

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #600

Topic 1

A company is planning to migrate a TCP-based application into the company's VPC. The application is publicly accessible on a nonstandard TCP port through a hardware appliance in the company's data center. This public endpoint can process up to 3 million requests per second with low latency. The company requires the same level of performance for the new public endpoint in AWS.

What should a solutions architect recommend to meet this requirement?

- A. Deploy a Network Load Balancer (NLB). Configure the NLB to be publicly accessible over the TCP port that the application requires. **Most Voted**
- B. Deploy an Application Load Balancer (ALB). Configure the ALB to be publicly accessible over the TCP port that the application requires.
- C. Deploy an Amazon CloudFront distribution that listens on the TCP port that the application requires. Use an Application Load Balancer as the origin.
- D. Deploy an Amazon API Gateway API that is configured with the TCP port that the application requires. Configure AWS Lambda functions with provisioned concurrency to process the requests.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Sugarbear_01 **Highly Voted** 1 year, 2 months ago

Selected Answer: A

Since the company requires the same level of performance for the new public endpoint in AWS.

A Network Load Balancer functions at the fourth layer of the Open Systems Interconnection (OSI) model. It can handle millions of requests per second. After the load balancer receives a connection request, it selects a target from the target group for the default rule. It attempts to open a TCP connection to the selected target on the port specified in the listener configuration.

Link;

<https://docs.aws.amazon.com/elasticloadbalancing/latest/network/introduction.html>

upvoted 11 times

TariqKipkemei Highly Voted 1 year ago

Selected Answer: A

TCP = NLB

upvoted 6 times

awsgeek75 Most Recent 11 months ago

Selected Answer: A

B: Is wrong as ALB is not going to help with TCP traffic

C: CloudFront is CDN. There is no content here

D: API Gateway is for HTTP web/API stuff, not custom TCP port applications

upvoted 3 times

taustin2 1 year, 2 months ago

Selected Answer: A

NLBs handle millions of requests per second. NLBs can handle general TCP traffic.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #601

Topic 1

A company runs its critical database on an Amazon RDS for PostgreSQL DB instance. The company wants to migrate to Amazon Aurora PostgreSQL with minimal downtime and data loss.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create a DB snapshot of the RDS for PostgreSQL DB instance to populate a new Aurora PostgreSQL DB cluster.
- B. Create an Aurora read replica of the RDS for PostgreSQL DB instance. Promote the Aurora read replicate to a new Aurora PostgreSQL DB cluster. **Most Voted**
- C. Use data import from Amazon S3 to migrate the database to an Aurora PostgreSQL DB cluster.
- D. Use the pg_dump utility to back up the RDS for PostgreSQL database. Restore the backup to a new Aurora PostgreSQL DB cluster.

Correct Answer: B*Community vote distribution***Comments****Guru4Cloud** **Highly Voted** 1 year, 8 months ago**Selected Answer: B**

The key reasons are:

Aurora read replicas allow setting up replication from RDS PostgreSQL to Aurora PostgreSQL with minimal downtime. Once replication is set up, the read replica can be promoted to a full standalone Aurora DB cluster with little to no downtime. This approach leverages AWS's managed replication between the source RDS PostgreSQL instance and Aurora. It avoids having to manually create backups and restore data.

Using DB snapshots or pg_dump backups requires manually restoring data which increases downtime and operational overhead.

Data import from S3 would require exporting, uploading and then importing data which adds overhead.

upvoted 7 times

JA2018 6 months, 2 weeks ago

can share the link? thanks

upvoted 1 times

Sugarbear_01 Highly Voted 1 year, 8 months ago

Answer [B]

There are five options for migrating data from your existing Amazon RDS for PostgreSQL database to an Amazon Aurora PostgreSQL-Compatible DB cluster.

- 1-Using a snapshot
- 2-Using an Aurora read replica
- 3-Using a pg_dump utility
- 4-Using logical replication
- 5-Using a data import from Amazon S3

(2-Using an Aurora read replica)

The Aurora read replica option minimizes downtime during a migration. Which is what the question demand so answer B; is the correct ;

<https://repost.aws/knowledge-center/aurora-postgresql-migrate-from-rds>

upvoted 5 times

Sugarbear_01 1 year, 8 months ago

Using (4 - using logical replication) RDS for PostgreSQL and Aurora PostgreSQL instance to migrate data off minimal downtime. But is not part of the option in the answer. Which makes answer B the best solution.

upvoted 2 times

Firdous586 Most Recent 1 year, 4 months ago

B is correct as the question says least down time and data loss

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

"Use an RDS for PostgreSQL DB instance as the basis for a new Aurora PostgreSQL DB cluster by using an Aurora read replica. The Aurora read replica is available for migrating only within the same AWS Region and account. The Aurora read replica option minimizes downtime during a migration. You can promote the new cluster when you have zero (0) replication lag between the primary RDS instance and the Aurora read replica."

<https://repost.aws/knowledge-center/aurora-postgresql-migrate-from-rds>

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: B

Not A: Would work but have some (though minor) downtime

B: "The Aurora read replica option minimizes downtime during a migration"

Not C: "If your data is stored using Amazon Simple Storage Service (Amazon S3)" ... in this case it is not

Not D: "If ... you don't have downtime considerations, you can use this option"

<https://repost.aws/knowledge-center/aurora-postgresql-migrate-from-rds>

upvoted 5 times

Cyberkayu 1 year, 5 months ago

Selected Answer: B

ACD will have delta changes issue. Which means, RDS snapshot/export at 2pm, upload/import the table into Aurora, configure and populated completed by 6pm. This created a 4-hour gap of delta changes

upvoted 2 times

aws94 1 year, 5 months ago

Selected Answer: A

please focus, we have RDS not Aurora, I don't know how you vote to create an Aurora read replica to migrate an RDS to Aurora.

upvoted 1 times

pentium75 1 year, 5 months ago

I thought that too but B is correct: <https://repost.aws/knowledge-center/aurora-postgresql-migrate-from-rds>

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

LEAST operational overhead = read replica
upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: B

A,B,C are all valid option.
But B: The Aurora read replica option minimizes downtime during a migration.
upvoted 2 times

thanhnv142 1 year, 7 months ago

B is correct guys. Lets see what we got here:
C and D is not correct of course. We have to consider A and B.
A: migration using a snapshot: this would, of course, introduce heavy data loss and down time
B: migration using read replica: nearly no dataloss and downtime.
upvoted 4 times

RRya 1 year, 7 months ago

Selected Answer: A

RDS PostgreSQL to Aurora PostgreSQL:
• Option 1: DB Snapshots from RDS PostgreSQL restored as PostgreSQL Aurora DB
• Option 2: Create an Aurora Read Replica from your RDS PostgreSQL, and when the replication lag is 0, promote it as its own DB cluster (can take time and cost \$)
upvoted 1 times

pentium75 1 year, 5 months ago

"The Aurora read replica option minimizes downtime during a migration"
<https://repost.aws/knowledge-center/aurora-postgresql-migrate-from-rds>
upvoted 2 times

Jay2k23 1 year, 8 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/AuroraPostgreSQL.Migrating.html>
upvoted 1 times

taustin2 1 year, 8 months ago

Selected Answer: B

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/AuroraPostgreSQL.Migrating.html>
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #602

Topic 1

A company's infrastructure consists of hundreds of Amazon EC2 instances that use Amazon Elastic Block Store (Amazon EBS) storage. A solutions architect must ensure that every EC2 instance can be recovered after a disaster.

What should the solutions architect do to meet this requirement with the LEAST amount of effort?

- A. Take a snapshot of the EBS storage that is attached to each EC2 instance. Create an AWS CloudFormation template to launch new EC2 instances from the EBS storage.
- B. Take a snapshot of the EBS storage that is attached to each EC2 instance. Use AWS Elastic Beanstalk to set the environment based on the EC2 template and attach the EBS storage.
- C. Use AWS Backup to set up a backup plan for the entire group of EC2 instances. Use the AWS Backup API or the AWS CLI to speed up the restore process for multiple EC2 instances. **Most Voted**
- D. Create an AWS Lambda function to take a snapshot of the EBS storage that is attached to each EC2 instance and copy the Amazon Machine Images (AMIs). Create another Lambda function to perform the restores with the copied AMIs and attach the EBS storage.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Guru4Cloud** **Highly Voted** 8 months, 2 weeks ago**Selected Answer: C**

The key reasons are:

AWS Backup automates backup of resources like EBS volumes. It allows defining backup policies for groups of resources. This removes the need to manually create backups for each resource.

The AWS Backup API and CLI allow programmatic control of backup plans and restores. This enables restoring hundreds of EC2 instances programmatically after a disaster instead of manually.

AWS Backup handles cleanup of old backups based on policies to minimize storage costs.

upvoted 13 times

LeonSauveterre **Most Recent** 6 months ago**Selected Answer: C**

A - Manually managing templates for hundreds of EC2 instances is TOO MUCH WORK.

B - Beanstalk is designed for managing application environments, not backups.

C - AWS Backup API or CLI simplifies bulk restore operations. YES.

D - Never create a lambda function when there's already a tool to do that and you need minimal effort. This is true for almost all the questions in this list.

upvoted 1 times

TariqKipkemei 6 months, 1 week ago

Selected Answer: C

LEAST amount of effort = AWS Backup

upvoted 3 times

Chiquitabandita 6 months, 3 weeks ago

for the question, I would choose C as well, AWS Backup of the EC2, but design, why would anything of importance be on the EC2 that would need to be restored? Shouldn't any critical or important data be on the EBS volumes in this example or similar location?

upvoted 2 times

LeonSauveterre 6 months ago

There are actually scenarios where an EC2 instance itself might still need restoration, even when EBS is being used for storage.

1. Some workloads may have specific configurations or dependencies baked into the EC2 instance itself, like custom softwares and OS-level configs.
2. Some applications may be running on EC2 relies on temporary but critical data stored on the instance's root volume.
3. Some stateful applications might store temporary state information locally on the EC2 instance. If this state isn't externalized (like DynamoDB, Redis, or EFS), then recovery is needed.

upvoted 1 times

taustin2 8 months, 2 weeks ago

Selected Answer: C

Going with Backup. Can restore programmatically using Backup API.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #603

Topic 1

A company recently migrated to the AWS Cloud. The company wants a serverless solution for large-scale parallel on-demand processing of a semistructured dataset. The data consists of logs, media files, sales transactions, and IoT sensor data that is stored in Amazon S3. The company wants the solution to process thousands of items in the dataset in parallel.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Use the AWS Step Functions Map state in Inline mode to process the data in parallel.
- B. Use the AWS Step Functions Map state in Distributed mode to process the data in parallel. **Most Voted**
- C. Use AWS Glue to process the data in parallel.
- D. Use several AWS Lambda functions to process the data in parallel.

Correct Answer: B*Community vote distribution***Comments****Guru4Cloud** **Highly Voted** 1 year, 8 months ago**Selected Answer: B**

AWS Step Functions allows you to orchestrate and scale distributed processing using the Map state. The Map state can process items in a large dataset in parallel by distributing the work across multiple resources.

Using the Map state in Distributed mode will automatically handle the parallel processing and scaling. Step Functions will add more workers to process the data as needed.

Step Functions is serverless so there are no servers to manage. It will scale up and down automatically based on demand.

upvoted 9 times

FlyingHawk **Most Recent** 4 months, 2 weeks ago**Selected Answer: B**

<https://docs.aws.amazon.com/step-functions/latest/dg/sample-dist-map-s3data-process.html>

With Step Functions, you can orchestrate large-scale parallel workloads to perform tasks, such as on-demand processing of semi-structured data. These parallel workloads let you concurrently process large-scale data sources stored in Amazon S3. For example, you might process a single JSON or CSV file that contains large amounts of data. Or you might process a large set of Amazon S3 objects.

upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

A - AWS Step Function Map State in inline mode accepts only a JSON array as input and supports up to 40 concurrent iterations.

upvoted 1 times

Sandy1254 11 months ago

Selected Answer: B

<https://docs.aws.amazon.com/step-functions/latest/dg/use-dist-map-orchestrate-large-scale-parallel-workloads.html>

upvoted 2 times

bogdannb 11 months, 2 weeks ago

Selected Answer: C

Using step functions will be overkill from my point of view. I would use Glue, it's serverless and purposely designed for such use case

upvoted 1 times

Lin878 11 months, 3 weeks ago

Selected Answer: B

Simple - user Lambda / Complex - user Step Functions

upvoted 3 times

Lx016 1 year, 4 months ago

A Map in Inline mode can support concurrency of 40 parallel branches and execution history limits of 25,000 events or approximately 6,500 state transitions in a workflow. With the Distributed mode, you can run at concurrency of up to 10,000 parallel branches. So I believe if it has to process thousands of items in parallel Distributed Mode is more appropriate

upvoted 4 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

<https://aws.amazon.com/blogs/aws/step-functions-distributed-map-a-serverless-solution-for-large-scale-parallel-data-processing/>

<https://docs.aws.amazon.com/step-functions/latest/dg/sample-dist-map-s3data-process.html>

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

The Distributed Map has been optimized for Amazon S3, helping you more easily iterate over objects in an S3 bucket. With the Distributed mode, you can run at concurrency of up to 10,000 parallel branches.

<https://aws.amazon.com/step-functions/faqs/#:~:text=A%20Map%20in%20Inline%20mode,up%20to%2010%2C000%20parallel%20branches.>

upvoted 3 times

Sugarbear_01 1 year, 8 months ago

Selected Answer: B

<https://docs.aws.amazon.com/step-functions/latest/dg/concepts-orchestrate-large-scale-parallel-workloads.html>

upvoted 2 times

taustin2 1 year, 8 months ago

Selected Answer: B

With Step Functions, you can orchestrate large-scale parallel workloads to perform tasks, such as on-demand processing of semi-structured data. These parallel workloads let you concurrently process large-scale data sources stored in Amazon S3.

<https://docs.aws.amazon.com/step-functions/latest/dg/concepts-orchestrate-large-scale-parallel-workloads.html>

upvoted 3 times

Sugarbear_01 1 year, 8 months ago

After going through the link I confirmed the answer is B

upvoted 2 times

[Removed] 1 year, 8 months ago

Large Scale + Parallel = Distributed Step Function

<https://docs.aws.amazon.com/step-functions/latest/dg/concepts-inline-vs-distributed-map.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #604

Topic 1


A company will migrate 10 PB of data to Amazon S3 in 6 weeks. The current data center has a 500 Mbps uplink to the internet. Other on-premises applications share the uplink. The company can use 80% of the internet bandwidth for this one-time migration task.

Which solution will meet these requirements?

- A. Configure AWS DataSync to migrate the data to Amazon S3 and to automatically verify the data.
- B. Use rsync to transfer the data directly to Amazon S3.
- C. Use the AWS CLI and multiple copy processes to send the data directly to Amazon S3.
- D. Order multiple AWS Snowball devices. Copy the data to the devices. Send the devices to AWS to copy the data to Amazon S3. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

Cyberkayu **Highly Voted** 1 year, 5 months ago

7 Years, 5 Months, 3 Weeks, 5 Days required to transfer 10PB on 400 Mbps. Finger cross the upload don't drop or timeout on year 7.

upvoted 12 times

zits88 9 months, 3 weeks ago

As a data engineer, this comment made both laugh and shudder at the same time -- hits too close to home

upvoted 2 times

TariqKipkemei **Highly Voted** 1 year, 6 months ago

Selected Answer: D

PB = snowball

upvoted 6 times

Ravan **Most Recent** 1 year, 3 months ago

Selected Answer: D

To calculate the total time required in weeks, we can use the result we obtained earlier, which was approximately 6.26
×
1
0
10
 6.26×10
10
weeks.

So, the total time required to transfer 10 PB of data to Amazon S3, given a 500 Mbps uplink, would be approximately 6.26
×
1
0
10
 6.26×10
10
weeks. However, this is an extremely large value and not practically feasible.

It's important to note that the result obtained might not accurately reflect real-world scenarios due to various factors such as network limitations, bandwidth constraints, and other practical considerations. Additionally, this calculation assumes a constant transfer rate and does not consider potential optimizations or parallelization techniques that could be employed to expedite the data transfer process.
upvoted 3 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

10PB on 80% of 500Mbps (Megabits not Megabytes) will take 6.5 years. But for the sake of exam when you cannot use calculators etc, just use snowball for petabytes of transfer if it is an option!
upvoted 2 times

awsgeek75 1 year, 4 months ago

Answer is D! not A! Fiddly fingers!
upvoted 7 times

wsdasdasdqwdaw 1 year, 7 months ago

D, but even if you do not know, all 3 option (A,B and C) have the same nature (transfer via bandwidth) and we know that there is only one correct answer => D.
upvoted 4 times

iwannabeawsgod 1 year, 7 months ago

Selected Answer: D

snowball for sure
upvoted 3 times

joshik 1 year, 8 months ago

Selected Answer: D

1Gbps will roughly do 7 TB in 24 hours. This means 400Mbps will only do 3x42TB.
upvoted 3 times

Xin123 1 year, 8 months ago

D
1Gbps will roughly do 7 TB in 24 hours. This means 400Mbps will only do 3x42TB.
upvoted 3 times

Sugarbear_01 1 year, 8 months ago

Selected Answer: D

D
1Gbps will roughly do 7 TB in 24 hours. This means 400Mbps will only do 3x42TB.

upvoted 3 times

Devsin2000 1 year, 8 months ago

D

1Gbps will roughly do 7 TB in 24 hours. This means 400Mbps will only do 3x42TB.

upvoted 2 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: D

D. Order multiple AWS Snowball devices. Copy the data to the devices. Send the devices to AWS to copy the data to Amazon S3.

upvoted 2 times

taustin2 1 year, 8 months ago

Selected Answer: D

10 PB = It's Snowballs.

upvoted 5 times

kambarami 1 year, 8 months ago

Answer is DDDDD

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #605

Topic 1

A company has several on-premises Internet Small Computer Systems Interface (iSCSI) network storage servers. The company wants to reduce the number of these servers by moving to the AWS Cloud. A solutions architect must provide low-latency access to frequently used data and reduce the dependency on on-premises servers with a minimal number of infrastructure changes.

Which solution will meet these requirements?

- A. Deploy an Amazon S3 File Gateway.
- B. Deploy Amazon Elastic Block Store (Amazon EBS) storage with backups to Amazon S3.
- C. Deploy an AWS Storage Gateway volume gateway that is configured with stored volumes.
- D. Deploy an AWS Storage Gateway volume gateway that is configured with cached volumes. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 2 months ago

Selected Answer: D

The key reasons are:

The Storage Gateway volume gateway provides iSCSI block storage using cached volumes. This allows replacing the on-premises iSCSI servers with minimal changes.
 Cached volumes store frequently accessed data locally for low latency access, while storing less frequently accessed data in S3. This reduces the number of on-premises servers while still providing low latency access to hot data.
 EBS does not provide iSCSI support to replace the existing servers.
 S3 File Gateway is for file storage, not block storage.
 Stored volumes would store all data on-premises, not in S3.
 upvoted 9 times

awsgeek75 **Most Recent** 11 months ago

Selected Answer: D

Low latency = always look for cache or local storage.
 A: Doesn't address low latency

A: Doesn't address low latency

B: Don't think this is possible

CD are both low latency but D is better:

<https://aws.amazon.com/storagegateway/faqs/#:~:text=In%20the%20cached%20mode%2C%20your,asynchronously%20backed%20up%20to%20AWS.>

upvoted 3 times

TariqKipkemei 1 year ago

Selected Answer: D

low-latency access to frequently used data = cached volumes

upvoted 4 times

Sugarbear_01 1 year, 2 months ago

Answer D

Here is the link ;

<https://docs.aws.amazon.com/storagegateway/latest/vgw/WhatIsStorageGateway.html>

upvoted 2 times

taustin2 1 year, 2 months ago

Selected Answer: D

ISCI=Volume Gateway.

low-latency access to frequently used data = cached volumes

upvoted 4 times

[Removed] 1 year, 2 months ago

"low-latency access to FREQUENTLY used data" = Cached AWS Storage Gateway volumes

upvoted 2 times

nnecode 1 year, 2 months ago

Selected Answer: D

An AWS Storage Gateway volume gateway is a hybrid storage solution that connects your on-premises applications to your cloud storage. It provides low-latency access to frequently used data while storing your entire dataset in the cloud.

When you configure an AWS Storage Gateway volume gateway with cached volumes, the gateway stores a copy of frequently accessed data locally. This allows you to provide low-latency access to your frequently accessed data while reducing your dependency on on-premises servers.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #606

Topic 1

A solutions architect is designing an application that will allow business users to upload objects to Amazon S3. The solution needs to maximize object durability. Objects also must be readily available at any time and for any length of time. Users will access objects frequently within the first 30 days after the objects are uploaded, but users are much less likely to access objects that are older than 30 days.

Which solution meets these requirements MOST cost-effectively?

- A. Store all the objects in S3 Standard with an S3 Lifecycle rule to transition the objects to S3 Glacier after 30 days.
- B. Store all the objects in S3 Standard with an S3 Lifecycle rule to transition the objects to S3 Standard-Infrequent Access (S3 Standard-IA) after 30 days. **Most Voted**
- C. Store all the objects in S3 Standard with an S3 Lifecycle rule to transition the objects to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 30 days.
- D. Store all the objects in S3 Intelligent-Tiering with an S3 Lifecycle rule to transition the objects to S3 Standard-Infrequent Access (S3 Standard-IA) after 30 days.

Correct Answer: B*Community vote distribution***Comments****Sugarbear_01** **Highly Voted** 1 year, 8 months ago**Selected Answer: B**

Minimum Days for Transition to S3 Standard-IA or S3 One Zone-IA

Before you transition objects to S3 Standard-IA or S3 One Zone-IA, you must store them for at least 30 days in Amazon S3. For example, you cannot create a Lifecycle rule to transition objects to the S3 Standard-IA storage class one day after you create them. Amazon S3 doesn't support this transition within the first 30 days because newer objects are often accessed more frequently or deleted sooner than is suitable for S3 Standard-IA or S3 One Zone-IA storage.

Similarly, if you are transitioning noncurrent objects (in versioned buckets), you can transition only objects that are at least 30 days noncurrent to S3 Standard-IA or S3 One Zone-IA storage.

copy, availability, durability, consistency, latency, availability, may also transition general considerations...

upvoted 5 times

abriggy Most Recent 10 months, 2 weeks ago

Answer is B.

C is wrong because one-zone doesn't maximize durability, it compromises it.

upvoted 3 times

MatAlves 8 months, 3 weeks ago

No, both Standard-IA and OZ-IA have SAME durability. What differs it the availability.

Durability is calculated based on how many times the data is copied and made redundant, inherently by the service itself. But that redundancy can be across one building (AZ) or multiple availability zones.

upvoted 2 times

TheLaPlanta 1 year, 2 months ago

Selected Answer: C

I believe it's C. The following link mentions One Zone-IA offers 99.999999999% durability. Questions says nothing about HA

upvoted 4 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

B
Intelligent tiering will automatically transition to S3 One Zone-IA which is not needed for durability.

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

'Objects also must be readily available at any time and for any length of time'...definitely option B.

upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: B

B is correct

upvoted 2 times

thanhnv142 1 year, 7 months ago

B is correct
C is not correct because data must be durable. C is only for data that can be regenerated.

upvoted 4 times

Xin123 1 year, 8 months ago

Selected Answer: B

Durability. Available any time for any duration => B

upvoted 2 times

Devsin2000 1 year, 8 months ago

A
S3 Glacier is most cost effective

upvoted 4 times

awsgeek75 1 year, 4 months ago

Between A & B, this is the tie-breaker:
"Objects also must be readily available at any time and for any length of time"

While Glacier IS more cost effective but it won't make the objects readily available at any time for any duration.... this is only possible with IA.

upvoted 2 times

taustin2 1 year, 8 months ago

Selected Answer: B

B meets the requirements. No need for intelligent Tiering because of 30 days.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #607

Topic 1

A company has migrated a two-tier application from its on-premises data center to the AWS Cloud. The data tier is a Multi-AZ deployment of Amazon RDS for Oracle with 12 TB of General Purpose SSD Amazon Elastic Block Store (Amazon EBS) storage. The application is designed to process and store documents in the database as binary large objects (blobs) with an average document size of 6 MB.

The database size has grown over time, reducing the performance and increasing the cost of storage. The company must improve the database performance and needs a solution that is highly available and resilient.

Which solution will meet these requirements MOST cost-effectively?

- A. Reduce the RDS DB instance size. Increase the storage capacity to 24 TiB. Change the storage type to Magnetic.
- B. Increase the RDS DB instance size. Increase the storage capacity to 24 TiB. Change the storage type to Provisioned IOPS.
- C. Create an Amazon S3 bucket. Update the application to store documents in the S3 bucket. Store the object metadata in the existing database. **Most Voted**
- D. Create an Amazon DynamoDB table. Update the application to use DynamoDB. Use AWS Database Migration Service (AWS DMS) to migrate data from the Oracle database to DynamoDB.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**awsgeek75** **Highly Voted** 11 months ago**Selected Answer: C**

When using BLOB, always try to pick a solution with S3.

upvoted 5 times

ferdzcrux **Most Recent** 10 months, 3 weeks ago

process and store documents as objects. S3 is known for object storage.

upvoted 4 times

TariqKipkemei 1 year ago

Selected Answer: C

MOST cost-effectively = store the objects in S3, and object metadata in the existing DB.

upvoted 3 times

taustin2 1 year, 2 months ago

DynamoDB's limit on the size of each record is 400KB, so D is wrong.

upvoted 3 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: C

C. Create an Amazon S3 bucket. Update the application to store documents in the S3 bucket. Store the object metadata in the existing database.

upvoted 4 times

taustin2 1 year, 2 months ago

Selected Answer: C

Storing the blobs in the db is more expensive than s3 with references in the db.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #608

Topic 1

A company has an application that serves clients that are deployed in more than 20,000 retail storefront locations around the world. The application consists of backend web services that are exposed over HTTPS on port 443. The application is hosted on Amazon EC2 instances behind an Application Load Balancer (ALB). The retail locations communicate with the web application over the public internet. The company allows each retail location to register the IP address that the retail location has been allocated by its local ISP.

The company's security team recommends to increase the security of the application endpoint by restricting access to only the IP addresses registered by the retail locations.

What should a solutions architect do to meet these requirements?

A. Associate an AWS WAF web ACL with the ALB. Use IP rule sets on the ALB to filter traffic. Update the IP addresses in the rule to include the registered IP addresses. **Most Voted**


B. Deploy AWS Firewall Manager to manage the ALB. Configure firewall rules to restrict traffic to the ALB. Modify the firewall rules to include the registered IP addresses.

C. Store the IP addresses in an Amazon DynamoDB table. Configure an AWS Lambda authorization function on the ALB to validate that incoming requests are from the registered IP addresses.

D. Configure the network ACL on the subnet that contains the public interface of the ALB. Update the ingress rules on the network ACL with entries for each of the registered IP addresses.

Correct Answer: A

Community vote distribution

**Comments**

pentium75 **Highly Voted** 11 months, 1 week ago

Selected Answer: A

WAF, you can have 100 "rule sets" per account, each with up to 10,000 IP addresses.

<https://docs.aws.amazon.com/waf/latest/developerguide/limits.html>

upvoted 8 times

Karls Most Recent 8 months ago

Selected Answer: C

AWS Lambda and DynamoDB to dynamically manage and validate incoming requests based on registered IP addresses.

<https://docs.aws.amazon.com/lambda/latest/dg/services-alb.html>

upvoted 1 times

ferdzacruz 10 months, 3 weeks ago

web services and HTTPS = WAF

upvoted 4 times

awsgeek75 11 months ago

Selected Answer: A

B: Looks like an incomplete solution for something different

C: Not workable as Lambda for IP filtering means you have already allowed the request to pass through

D NACL with entries for each registered IP is not possible.

upvoted 3 times

TariqKipkemei 1 year ago

Selected Answer: A

endpoint restriction by IP addresses = AWS WAF

upvoted 4 times

Passexam4sure_com 1 year, 1 month ago

Selected Answer: A

Associate an AWS WAF web ACL with the ALB. Use IP rule sets on the ALB to filter traffic. Update the IP addresses in the rule to include the registered IP addresses.

upvoted 4 times

Sugarbear_01 1 year, 2 months ago

Selected Answer: A

AWS WAF cannot be directly associated with a Web Application. But, can only be associated with Application Load Balancer, CloudFront and API Gateway.

upvoted 4 times

taustin2 1 year, 2 months ago

Selected Answer: C

Changing answer to C because of "20000" IP addresses. Use Lambda with ALB.

upvoted 3 times

pentium75 11 months, 1 week ago

WAF allows 100 rule sets, each with up to 10,000 IP addresses, per account.

upvoted 2 times

potomac 1 year, 1 month ago

10,000 IP addresses

For the latest version of AWS WAF, see AWS WAF. If you want to allow or block web requests based on the IP addresses that the requests originate from, create one or more IP match conditions. An IP match condition lists up to 10,000 IP addresses or IP address ranges that your requests originate from.

upvoted 2 times

bsbs1234 1 year, 1 month ago

I will choose this answer if it is API Gateway. But I cannot figure out how to do lambda authentication on ALB. I will go A

upvoted 2 times

taustin2 1 year, 1 month ago

You are right. I don't know of a way to use Lambda with ALB in this way. Answer is A.

upvoted 2 times

upvoted 2 times

potomac 1 year, 1 month ago

ALB invokes Lambda function, sending the incoming data in JSON format. Lambda function performs task, returns HTTP response to ALB.

upvoted 1 times

potomac 1 year, 1 month ago

WAF seems still better

upvoted 3 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: A

A. Associate an AWS WAF web ACL with the ALB. Use IP rule sets on the ALB to filter traffic. Update the IP addresses in the rule to include the registered IP addresses.

upvoted 3 times

taustin2 1 year, 2 months ago

Selected Answer: A

WAF meets the requirements.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #609

Topic 1

A company is building a data analysis platform on AWS by using AWS Lake Formation. The platform will ingest data from different sources such as Amazon S3 and Amazon RDS. The company needs a secure solution to prevent access to portions of the data that contain sensitive information.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an IAM role that includes permissions to access Lake Formation tables.
- B. Create data filters to implement row-level security and cell-level security. **Most Voted**
- C. Create an AWS Lambda function that removes sensitive information before Lake Formation ingests the data.
- D. Create an AWS Lambda function that periodically queries and removes sensitive information from Lake Formation tables.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**Guru4Cloud** **Highly Voted** 1 year, 8 months ago**Selected Answer: B**

The key reasons are:

Lake Formation data filters allow restricting access to rows or cells in data tables based on conditions. This allows preventing access to sensitive data.

Data filters are implemented within Lake Formation and do not require additional coding or Lambda functions.

Lambda functions to pre-process data or purge tables would require ongoing development and maintenance.

IAM roles only provide user-level permissions, not row or cell level security.

Data filters give granular access control over Lake Formation data with minimal configuration, avoiding complex custom code.

upvoted 10 times

awsgeek75 1 year, 4 months ago

<https://docs.aws.amazon.com/lake-formation/latest/dg/data-filters-about.html>

upvoted 2 times

emakid **Most Recent** 11 months, 2 weeks ago

Chandra 11 months, 2 weeks ago

Selected Answer: B

B. Create data filters to implement row-level security and cell-level security.

Explanation:

Row-Level and Cell-Level Security: AWS Lake Formation provides built-in support for row-level and cell-level security. By using data filters, you can define policies that control access to specific rows and cells within your tables. This allows you to restrict access to sensitive information without needing to manually filter or remove data.

Least Operational Overhead: This solution leverages built-in Lake Formation capabilities, reducing the need for additional infrastructure or custom code. Once the data filters are set up, they automatically enforce the security policies, minimizing ongoing operational overhead.

upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: B

As it said "prevent access to portions of the data that contain sensitive information", not the access to S3, so data filter is enough

upvoted 3 times

wizcloudifa 1 year, 1 month ago

Selected Answer: B

Focus on the exact wordings: "to prevent access to portions of the data that contain sensitive information."

Only option B restricts the platform to access sensitive data, option A restrict users to restrict access that doesn't serve the req here, C and D are talking about removing the sensitive data which is not the ask here

upvoted 2 times

ferdzcruz 1 year, 4 months ago

portions of the data that contain sensitive information = Filtered data.

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

A is possible but it does not secure the data properly and only provides table level access control (if any).

CD are too much overhead

B is exactly for this purpose and is a built-in feature of Lake formation

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: B

<https://docs.aws.amazon.com/lake-formation/latest/dg/data-filters-about.html>

upvoted 3 times

taustin2 1 year, 8 months ago

Selected Answer: B

You can create data filters based on the values of columns in a Lake Formation table. Easy. Lowest operational overhead.

upvoted 3 times

nnecode 1 year, 8 months ago

Selected Answer: B

The best solution to meet the requirements with the least operational overhead is to create data filters to implement row-level security and cell-level security.

Data filters are a feature of Lake Formation that allow you to restrict access to data based on row and column values. This can be used to implement row-level security and cell-level security.

To implement row-level security, you would create a data filter that only allows users to access rows where the values in certain columns meet certain criteria. For example, you could create a data filter that only allows users to access rows where the value in the customer_id column matches the user's own customer ID.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #610

Topic 1


A company deploys Amazon EC2 instances that run in a VPC. The EC2 instances load source data into Amazon S3 buckets so that the data can be processed in the future. According to compliance laws, the data must not be transmitted over the public internet. Servers in the company's on-premises data center will consume the output from an application that runs on the EC2 instances.

Which solution will meet these requirements?

- A. Deploy an interface VPC endpoint for Amazon EC2. Create an AWS Site-to-Site VPN connection between the company and the VPC.
- B. Deploy a gateway VPC endpoint for Amazon S3. Set up an AWS Direct Connect connection between the on-premises network and the VPC. **Most Voted**
- C. Set up an AWS Transit Gateway connection from the VPC to the S3 buckets. Create an AWS Site-to-Site VPN connection between the company and the VPC.
- D. Set up proxy EC2 instances that have routes to NAT gateways. Configure the proxy EC2 instances to fetch S3 data and feed the application instances.

Correct Answer: B

Community vote distribution



Comments

taustin2 **Highly Voted** 1 year, 8 months ago

Selected Answer: B

Gateway VPC Endpoint = no internet to access S3. Direct Connect = secure access to VPC.
upvoted 11 times

awsgeek75 **Highly Voted** 1 year, 4 months ago

Selected Answer: B

No public internet != encrypted public internet (VPN)
Direct connect is the only option.
upvoted 6 times

zdi561 Most Recent 4 months ago

Selected Answer: A

While not as ideal for high-traffic scenarios, a properly configured Site-to-Site VPN can also create a private connection between your on-premise network and AWS. Gateway endpoint is not working for on-premise
upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: B

Deploy a gateway VPC endpoint for Amazon S3 = so traffic between EC2 and S3 doesn't live AWS private network.

Set up an AWS Direct Connect connection between the on-premises network and the VPC = servers on-premises can consume the output from ec2 instances via private connection.
upvoted 2 times

OSHOAIB 1 year, 5 months ago

Selected Answer: B

A gateway VPC endpoint for Amazon S3 allows the EC2 instances within the VPC to access Amazon S3 buckets without using the public internet. The traffic between the VPC and S3 is routed within the AWS network.
AWS Direct Connect establishes a private connection between the on-premises data center and AWS infrastructure, avoiding data transfer over the public internet and ensuring compliance with the specified requirements. It provides a dedicated network link with higher bandwidth options and potentially more consistent network performance than internet-based connections. Whereas Option A uses Site-to-Site VPN connection which is secure. However it typically runs over the public internet, which would not meet the company's requirement of avoiding public internet data transit.
upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: B

I think the last sentence ("Servers in the company's on-premises data center will consume the output from an application that runs on the EC2 instances") refers to a different application. Purely from the wording, it does NOT seem to refer to the data 'loaded into S3 buckets so that it can be processed in the future' before. So the EC2 instances could write to S3, the on-premises servers can talk to the EC2 application, and data would not be transmitted over the public internet.

Not A: There's no such thing as a "VPC endpoint for Amazon EC2 (!)"
Not C: Transit Gateway is not for EC2->S3, VPN is over public internet
Not D: Would address only the first part and use public Internet
upvoted 2 times

wizcloudifa 1 year, 1 month ago

Interface endpoint is a thing, the only reason A is not true is because of the presence of site-to-site vpn which is essentially accessing public internet
upvoted 2 times

ale_brd_111 1 year, 5 months ago

Selected Answer: A

I would go for A, for two reasons:
1) "S3 gateway endpoints do not currently support access from resources in a different Region, different VPC, or from an on-premises (non-AWS) environment."
2) we tryna access an output from an application hosted in e2 instances and not to access the s3 stored data so ideally we should use Interface Endpoints for the applications running in ec2.
upvoted 3 times

MatAlves 8 months, 3 weeks ago

You forgot the traffic from EC2 to S3. Without the Gateway Endpoint, that would go via public internet.

1. Deploy a gateway VPC endpoint for Amazon S3 = so traffic between EC2 and S3 doesn't live AWS private network.
2. Set up an AWS Direct Connect connection between the on-premises network and the VPC = servers on-premises can consume the output from ec2 instances via private connection.
upvoted 2 times

pentium75 1 year, 5 months ago

Plus, in A you deploy a VPC endpoint "for EC2" (!) which doesn't exist
upvoted 3 times

elmyth 9 months ago

exists, check the docs, interface VPS endpoint != gateway VPC endpoint, they have different range of services
upvoted 2 times

pentium75 1 year, 5 months ago

"Data must not be transmitted over the public internet", as it would with A (VPN).
upvoted 3 times

ftaws 1 year, 5 months ago

I standhood answer is B, but why not A?
upvoted 1 times

pentium75 1 year, 5 months ago

there's no such things a 'VPC endpoint for EC2', and it uses public Internet
upvoted 2 times

achechen 1 year, 6 months ago

Selected Answer: A

<https://aws.amazon.com/blogs/architecture/choosing-your-vpc-endpoint-strategy-for-amazon-s3/> According to this document, " S3 gateway endpoints do not currently support access from resources in a different Region, different VPC, or from an on-premises (non-AWS) environment. However, if you're willing to manage a complex custom architecture, you can use proxies. In all those scenarios, where access is from resources external to VPC, S3 interface endpoints access S3 in a secure way." so, the answer is A.
upvoted 4 times

pentium75 1 year, 5 months ago

A uses a VPC endpoint "for Amazon EC2", not S3. Also it uses public Internet.
upvoted 3 times

elmyth 9 months ago

interface VPC endpoint works with PrivateLink, so it can be connected to huge amount of services, and to EC2. Gateway VPC endpoint can't work for on-prem
upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

data must not be transmitted over the public internet = gateway VPC endpoint for Amazon S3 and AWS Direct Connect connection between the on-premises network and the VPC.
upvoted 2 times

Guru4Cloud 1 year, 8 months ago

Selected Answer: B

Gateway VPC Endpoint = no internet to access S3. Direct Connect = secure access to VPC
I agree with you @taustin2- Happy Learning all
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #611

Topic 1

A company has an application with a REST-based interface that allows data to be received in near-real time from a third-party vendor. Once received, the application processes and stores the data for further analysis. The application is running on Amazon EC2 instances.

The third-party vendor has received many 503 Service Unavailable Errors when sending data to the application. When the data volume spikes, the compute capacity reaches its maximum limit and the application is unable to process all requests.

Which design should a solutions architect recommend to provide a more scalable solution?

- A. Use Amazon Kinesis Data Streams to ingest the data. Process the data using AWS Lambda functions. **Most Voted**
- B. Use Amazon API Gateway on top of the existing application. Create a usage plan with a quota limit for the third-party vendor.
- C. Use Amazon Simple Notification Service (Amazon SNS) to ingest the data. Put the EC2 instances in an Auto Scaling group behind an Application Load Balancer.
- D. Repackage the application as a container. Deploy the application using Amazon Elastic Container Service (Amazon ECS) using the EC2 launch type with an Auto Scaling group.

Correct Answer: A*Community vote distribution*

A (94%)

B (6%)

Comments**Guru4Cloud** **Highly Voted** 1 year, 2 months ago**Selected Answer: A**

The key reasons are:

Kinesis Data Streams provides an auto-scaling stream that can handle large amounts of streaming data ingestion and throughput. This removes the bottlenecks around receiving the data.

AWS Lambda can process and store the data in a scalable serverless manner, avoiding EC2 capacity limits.

API Gateway adds API management capabilities but does not improve the underlying scalability of the EC2 application.

SNS is for event publishing/notifications, not large scale data ingestion. ECS still relies on EC2 capacity.

upvoted 7 times

taustin2 Highly Voted 1 year, 2 months ago

Selected Answer: A

For near-real time data ingest and processing, Kinesis and Lambda are most scalable choice.

upvoted 5 times

zdi561 Most Recent 4 months ago

Selected Answer: B

A change the whole architect, B adding a API gateway to make the app more scalable not perfect as in A it may not be possible within budget, deadline, Yes, AWS API Gateway significantly contributes to making an application more scalable by handling a large volume of concurrent API requests, automatically managing traffic distribution, it can cache response and allow asyn call

upvoted 1 times

ferdzacruz 10 months, 3 weeks ago

A.

Kinesis Data Streams = near realtime and scalable

AWS Lambda functions = scalable

upvoted 3 times

TariqKipkemei 1 year ago

Selected Answer: A

more scalable solution? = serverless = Amazon Kinesis Data Streams and AWS Lambda functions

upvoted 3 times

wsdadsdqdaw 1 year, 1 month ago

Only A is pure serverless which means scale. A for sure.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #612

Topic 1

A company has an application that runs on Amazon EC2 instances in a private subnet. The application needs to process sensitive information from an Amazon S3 bucket. The application must not use the internet to connect to the S3 bucket.

Which solution will meet these requirements?

- A. Configure an internet gateway. Update the S3 bucket policy to allow access from the internet gateway. Update the application to use the new internet gateway.
- B. Configure a VPN connection. Update the S3 bucket policy to allow access from the VPN connection. Update the application to use the new VPN connection.
- C. Configure a NAT gateway. Update the S3 bucket policy to allow access from the NAT gateway. Update the application to use the new NAT gateway.
- D. Configure a VPC endpoint. Update the S3 bucket policy to allow access from the VPC endpoint. Update the application to use the new VPC endpoint. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 2 months ago

Selected Answer: D

The solution that will meet these requirements is to:

Configure a VPC endpoint for Amazon S3
Update the S3 bucket policy to allow access from the VPC endpoint
Update the application to use the new VPC endpoint
The key reasons are:

VPC endpoints allow private connectivity from VPCs to AWS services like S3 without using an internet gateway. The application can connect to S3 through the VPC endpoint while remaining in the private subnet, without internet access.
upvoted 8 times

ferdzcruz **Most Recent** 10 months, 3 weeks ago

D.

VPC endpoint = not internet, direct access from VPC to S3

upvoted 3 times

awsgeek75 11 months ago

Selected Answer: D

<https://docs.aws.amazon.com/whitepapers/latest/aws-privatelink/what-are-vpc-endpoints.html>

upvoted 3 times

achechen 1 year ago

Selected Answer: D

Answer is D

upvoted 4 times

TariqKipkemei 1 year ago

Selected Answer: D

application must not use the internet to connect to the S3 bucket = VPC endpoint

upvoted 4 times

taustin2 1 year, 2 months ago

Selected Answer: D

VPC Endpoint for S3.

upvoted 2 times

aleariva 1 year, 2 months ago

D is the correct...<https://docs.aws.amazon.com/whitepapers/latest/aws-privatelink/what-are-vpc-endpoints.html>

upvoted 1 times

awslearnerin2022 1 year, 2 months ago

Selected Answer: D

VPC endpoint enables communication between VPC subnet and S3 bucket.

upvoted 1 times

nnecode 1 year, 2 months ago

Selected Answer: D

A VPC endpoint is a managed endpoint in your VPC that is connected to a public AWS service. It provides a private connection between your VPC and the service, and it does not require an internet gateway or a NAT device.

Option A (internet gateway) would involve exposing the S3 bucket to the internet, which is not recommended for security reasons.

Option B (VPN connection) would require additional setup and would still involve traffic going over the internet.

Option C (NAT gateway) is used for outbound internet access from private subnets, not for accessing S3 without the internet.

upvoted 4 times