

C could work but is less managed and "configure the backend code to reference EC2 instance" makes no sense
D ECS+Linux+PHP is good managed combination when used with Fargate. S3 for static is well-architected. Multi-AZ ElastiCache for Redis is HA also. Good managed solution for all purposes.

upvoted 5 times

ferdzcruz Most Recent 10 months, 3 weeks ago

D. ECS + Fargate

Company wants to redesign the architecture = from Server to serverless, and managed by AWS .

upvoted 4 times

evelynsun 11 months, 3 weeks ago

Selected Answer: D

This solution meets the requirements because it uses AWS managed solutions for hosting the static content and the PHP application. It also uses Amazon ECS to run the PHP application in a highly available and scalable manner. The solution also uses Amazon ElastiCache for Redis to handle session information, which is highly available and scalable. The solution also uses Amazon CloudFront to provide a secure and reliable way to deliver the static content to users.

upvoted 4 times

TariqKipkemei 1 year ago

Selected Answer: D

Configure an Amazon CloudFront distribution with an Amazon S3 endpoint to an S3 bucket that is configured to host the static content. Configure an Application Load Balancer that targets an Amazon Elastic Container Service (Amazon ECS) service that runs AWS Fargate tasks for the PHP application. Configure the PHP application to use an Amazon ElastiCache for Redis cluster that runs in multiple Availability Zones.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #655

Topic 1

A company runs a web application on Amazon EC2 instances in an Auto Scaling group that has a target group. The company designed the application to work with session affinity (sticky sessions) for a better user experience.

The application must be available publicly over the internet as an endpoint. A WAF must be applied to the endpoint for additional security. Session affinity (sticky sessions) must be configured on the endpoint.

Which combination of steps will meet these requirements? (Choose two.)

- A. Create a public Network Load Balancer. Specify the application target group.
- B. Create a Gateway Load Balancer. Specify the application target group.

C. Create a public Application Load Balancer. Specify the application target group. **Most Voted**

D. Create a second target group. Add Elastic IP addresses to the EC2 instances.

E. Create a web ACL in AWS WAF. Associate the web ACL with the endpoint **Most Voted**

Correct Answer: CE

Community vote distribution



Comments

ferdzcruz 1 year, 4 months ago

CE.
C. application = ALB
E. WAF to endpoint
upvoted 4 times

awsgeek75 1 year, 4 months ago

Selected Answer: CE

NLB and GLB cannot handle sticky sessions. It's an application level concept (Cookies) so ALB works.
Elastic IP will negate sticky sessions and this combination won't work.
E give proper permissions to WAF
upvoted 4 times

Amal Kumar 7 months ago

tonyduivannngnia / months, 2 weeks ago

I agree with you CE is correct but your explanation is wrong. NLB, ALB and CLB can handle sticky sessions. But the WAF is just only working with ALB, so ALB is correct.

upvoted 1 times

LeonSauveterre 5 months, 3 weeks ago

After a little digging, turns out NLBs only offer IP-based stickiness, which is not ideal for web applications where many users might share the same public IP address (due to NAT oh well). Cookie-based stickiness offered by ALBs is more appropriate.

PS. Sticky sessions are not supported if you are using TLS termination on NLB. If you need TLS termination and sticky sessions, you should use ALB as well.

upvoted 1 times

Mikado211 1 year, 5 months ago

Selected Answer: CE

- Make it accessible from the web + sticky session == Public ALB
- Additional security == web ACL in WAF (and integrate the web ACL to the ALB)

upvoted 2 times

ZZZ_Sleep 1 year, 5 months ago

Selected Answer: CE

session affinity (sticky sessions) = Application Load Balancer

WAF must be applied to the endpoint for additional security = web ACL in WAF

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: CE

Session Affinity = Application Load Balancer

Create a public Application Load Balancer. Specify the application target group then create a web ACL in AWS WAF. Associate the web ACL with the ALB endpoint.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #656

Topic 1


A company runs a website that stores images of historical events. Website users need the ability to search and view images based on the year that the event in the image occurred. On average, users request each image only once or twice a year. The company wants a highly available solution to store and deliver the images to users.

Which solution will meet these requirements MOST cost-effectively?

- A. Store images in Amazon Elastic Block Store (Amazon EBS). Use a web server that runs on Amazon EC2.
- B. Store images in Amazon Elastic File System (Amazon EFS). Use a web server that runs on Amazon EC2.
- C. Store images in Amazon S3 Standard. Use S3 Standard to directly deliver images by using a static website.
- D. Store images in Amazon S3 Standard-Infrequent Access (S3 Standard-IA). Use S3 Standard-IA to directly deliver images by using a static website. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

chikuwan **Highly Voted** 1 year, 6 months ago

Selected Answer: D

users request each image only once or twice a year
So the answer is D
upvoted 10 times

Abbas_Abi_AWS **Most Recent** 10 months, 1 week ago

Selected Answer: D

D is correct
upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: D

users request each image only once or twice a year, so infrequent is enough and cheaper
upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: D

"On average, users request each image only once or twice a year."
S3 Infrequent Access is more than enough for this.

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: D

Say, you have 1 TB of files that you access twice a year. Yearly cost:
C, S3 Standard: 276 USD for storage, free retrieval = 276 USD
D, S3 Standard-IA: 138 USD for storage, 20 € for retrieval = 158 USD

upvoted 2 times

Kumar05162 1 year, 5 months ago

Option D: Store images in Amazon S3 Standard-Infrequent Access (S3 Standard-IA). Use S3 Standard-IA to directly deliver images by using a static website.

S3 Standard-IA is designed specifically for infrequently accessed data, offering lower storage costs compared to S3 Standard while still providing the necessary durability and availability.

upvoted 2 times

ZZZ_Sleep 1 year, 5 months ago

Selected Answer: D

High Availability = excluded A (EBS)
cost-effective = excluded B (EFS)
only once or twice a year = S3 Standard-IA, excluded C (S3 Standard, frequent access)

Left D, answer

upvoted 4 times

LuADS 1 year, 6 months ago

Selected Answer: C

Suppose there are thousands or millions of users, each image should be recovered once or twice a year X total users... makes it more expensive than the standard class since the recovery price of Standard-IA is \$0.01 per GB + price of the requests which is also more expensive too.

upvoted 2 times

pentium75 1 year, 5 months ago

Not sure if you understood what IA class does. "Recovery price is 0.001 per GB", what's the issue with that if images are requested "only once or twice a year"?

Say, you have 1 TB of files that you access twice a year.
S3 Standard: 276 USD for storage, free retrieval = 276 USD
S3 Standard-IA: 138 USD for storage, 20 € for retrieval = 158 USD

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: D

MOST cost-effectively, request each image only once or twice a year= S3 Standard-Infrequent Access

upvoted 2 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

D
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/storage-class-intro.html>
Look at table

upvoted 2 times

achechen 1 year, 6 months ago

Selected Answer: D

if the images are accessed once or twice a year, then it is cheaper to use infrequent access tier
upvoted 4 times

aragornfsm 1 year, 6 months ago

I believe the correct answer is option D, but ChatGPT mentioned option C. I didn't understand. I'm curious about the actual correct answer.
upvoted 2 times

AndreiWebNet 1 year, 6 months ago

Might be the fact the a user is requesting to view a image once or twice a year but how many users are there ? :) that's why it points to C i think. I still think that the correct answer is D due to lack of information in the description
upvoted 1 times

pentium75 1 year, 5 months ago

"Users request each image only once or twice per year", this refers to all users together, it does not say "EACH user". In other words, every image is accessed once or twice a year.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #657

Topic 1

A company has multiple AWS accounts in an organization in AWS Organizations that different business units use. The company has multiple offices around the world. The company needs to update security group rules to allow new office CIDR ranges or to remove old CIDR ranges across the organization. The company wants to centralize the management of security group rules to minimize the administrative overhead that updating CIDR ranges requires.

Which solution will meet these requirements MOST cost-effectively?

A. Create VPC security groups in the organization's management account. Update the security groups when a CIDR range update is necessary.

B. Create a VPC customer managed prefix list that contains the list of CIDRs. Use AWS Resource Access Manager (AWS RAM) to share the prefix list across the organization. Use the prefix list in the security groups across the organization.

Most Voted

C. Create an AWS managed prefix list. Use an AWS Security Hub policy to enforce the security group update across the organization. Use an AWS Lambda function to update the prefix list automatically when the CIDR ranges change.

D. Create security groups in a central administrative AWS account. Create an AWS Firewall Manager common security group policy for the whole organization. Select the previously created security groups as primary groups in the policy.

Correct Answer: B*Community vote distribution*B (100%)**Comments****TariqKipkemei** **Highly Voted** 1 year, 6 months ago**Selected Answer: B**

A managed prefix list is a set of one or more CIDR blocks. You can use prefix lists to make it easier to configure and maintain your security groups and route tables. You can create a prefix list from the IP addresses that you frequently use, and reference them as a set in security group rules and routes instead of referencing them individually. If you scale your network and need to allow traffic from another CIDR block, you can update the relevant prefix list and all security groups that use the prefix list are updated. You can also use managed prefix lists with other AWS accounts using Resource Access Manager (RAM).

<https://docs.aws.amazon.com/vpc/latest/userguide/managed-prefix-lists.html#:~:text=A-,managed%20prefix,-list%20is%20a>
upvoted 8 times

awsgeek75 Highly Voted 1 year, 4 months ago

Such a badly worded question:

"The company has multiple offices around the world. The company needs to update security group rules to allow new office CIDR ranges or to remove old CIDR ranges across the organization."

Are the CIDR groups associated to offices? That will be illogical. I think it should be VPC and not offices.

upvoted 5 times

Gape4 Most Recent 11 months, 1 week ago

Selected Answer: B

I will go for B

upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: B

prefix list for CIDR blocks

upvoted 2 times

avdxeqtr 1 year, 4 months ago

Selected Answer: B

<https://docs.aws.amazon.com/vpc/latest/userguide/managed-prefix-lists.html>

upvoted 2 times

ale_brd_111 1 year, 5 months ago

Selected Answer: B

Answer is B

upvoted 2 times

achechen 1 year, 6 months ago

Selected Answer: B

looks like B is the answer. Reference: <https://docs.aws.amazon.com/vpc/latest/userguide/managed-prefix-lists.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #658

Topic 1

A company uses an on-premises network-attached storage (NAS) system to provide file shares to its high performance computing (HPC) workloads. The company wants to migrate its latency-sensitive HPC workloads and its storage to the AWS Cloud. The company must be able to provide NFS and SMB multi-protocol access from the file system.

Which solution will meet these requirements with the LEAST latency? (Choose two.)

A. Deploy compute optimized EC2 instances into a cluster placement group. **Most Voted**

B. Deploy compute optimized EC2 instances into a partition placement group.


C. Attach the EC2 instances to an Amazon FSx for Lustre file system.

D. Attach the EC2 instances to an Amazon FSx for OpenZFS file system.

E. Attach the EC2 instances to an Amazon FSx for NetApp ONTAP file system. **Most Voted**

Correct Answer: AE

Community vote distribution

**Comments**

lucasbg **Highly Voted** 1 year, 6 months ago

Selected Answer: AE

You talked about smb and nfs, you talked fsx netapp ontap

C is wrong because lustre is a POSIX fs
upvoted 13 times

awsgeek75 **Highly Voted** 1 year, 4 months ago

Selected Answer: AE

A Because HPC equivalent in AWS is EC2. Cluster placement for low-latency:
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>
E: ONTAP gives NFS and SMB which is required
AE is correct
B does not solve low latency requirements
C No support for NFS and SMB
D OpenZFS is not required

upvoted 7 times

awsgeek75 1 year, 4 months ago

<https://aws.amazon.com/fsx/netapp-ontap/features/>

Amazon FSx for NetApp ONTAP provides access to shared file storage over all versions of the Network File System (NFS) and Server Message Block (SMB) protocols, and also supports multi-protocol access (i.e. concurrent NFS and SMB access) to the same data. As a result, you can access Amazon FSx for NetApp ONTAP from virtually any Linux, Windows, or macOS client.

upvoted 2 times

MatAlves Most Recent 8 months, 3 weeks ago

Selected Answer: AE

I got baited into quickly going for Lustre after reading, but forgot it doesn't not support NFS/SMB.

upvoted 4 times

tsdsmth 1 year, 4 months ago

Amazon FSx for Lustre does not support SMB. So it's A,E

upvoted 6 times

1rob 1 year, 5 months ago

Selected Answer: AD

A because cluster placement group means low latency, and D because OpenZFS has less latency compared to FSx for NetApp ONTAP. See <https://aws.amazon.com/fsx/when-to-choose-fsx/>
FSx for OpenZFS can handle SMB and NFS.

Despite that for onprem NAS appliances the aws Recommended Amazon FSx file system would be FSx for NetApp ONTAP, I still choose FSx for OpenZFS for the lower latency.

upvoted 2 times

ZZZ_Sleep 1 year, 5 months ago

Selected Answer: AE

LEAST latency = cluster placement group

Amazon FSx for Lustre = SMB

Amazon FSx for OpenZFS = NFS

Amazon FSx for NetApp ONTAP = NFS, SMB, iSCSI

So, answer are A and E

upvoted 7 times

Sumith4112 1 year, 6 months ago

Selected Answer: AE

A because cluster placement group means low latency.

E

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: AE

HPC, NFS, SMB = FSx for NetApp ONTAP file system

HPC, latency-sensitive = cluster placement group

upvoted 4 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

AE

<https://aws.amazon.com/fsx/when-to-choose-fsx/>

upvoted 2 times

achechen 1 year, 6 months ago

Selected Answer: AE

I don't think FSx for Lustre supports SMB. At least I could not find anything in the documentation. However, FSx for ONTAP delivers NFS and SMB support

supports NFS and SMB support.

upvoted 4 times

chikuwan 1 year, 6 months ago

Selected Answer: AE

<https://aws.amazon.com/jp/fsx/lustre/features/>

upvoted 4 times

reika1914 1 year, 6 months ago

Selected Answer: AC

To meet the requirements of migrating latency-sensitive HPC workloads with multi-protocol access (NFS and SMB) to AWS with minimal latency, the following solutions would be the most appropriate:

A. Deploy compute optimized EC2 instances into a cluster placement group.

C. Attach the EC2 instances to an Amazon FSx for Lustre file system.

upvoted 2 times

pentium75 1 year, 5 months ago

FSx for Lustre provides Lustre, not SMB and not NFS

upvoted 2 times

Chiquitabandita 1 year, 6 months ago

Selected Answer: AE

[https://aws.amazon.com/fsx/netapp-ontap/features/#:~:text=Amazon%20FSx%20for%20NetApp%20ONTAP%20provides%20access%20to%20shared%20file,access\)%20to%20the%20same%20data.](https://aws.amazon.com/fsx/netapp-ontap/features/#:~:text=Amazon%20FSx%20for%20NetApp%20ONTAP%20provides%20access%20to%20shared%20file,access)%20to%20the%20same%20data.) "Amazon FSx for NetApp ONTAP provides access to shared file storage over all versions of the Network File System (NFS) and Server Message Block (SMB) protocols, and also supports multi-protocol access (i.e. concurrent NFS and SMB access) to the same data."

upvoted 5 times

LemonGremlin 1 year, 6 months ago

Selected Answer: AC

Option A: A cluster placement group provides low-latency and high-bandwidth connectivity between instances. This is particularly beneficial for high-performance computing workloads that are latency-sensitive. Instances within a cluster placement group are placed in close proximity to each other within the same Availability Zone.

Option C: Amazon FSx for Lustre is a high-performance file system optimized for fast access to data. It is well-suited for high-performance computing workloads. It provides low-latency access to data and supports the NFS protocol.

upvoted 3 times

t0nx 1 year, 6 months ago

Thank you

upvoted 1 times

1rob 1 year, 5 months ago

FSx for Lustre is not about NFS or SMB. You will need a linux instance. First install the open-source Lustre client on that instance. Once it's installed, you can mount your file system using standard Linux commands. So C is not correct here because NFS and SMB support is required.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #659

Topic 1

A company is relocating its data center and wants to securely transfer 50 TB of data to AWS within 2 weeks. The existing data center has a Site-to-Site VPN connection to AWS that is 90% utilized.

Which AWS service should a solutions architect use to meet these requirements?

A. AWS DataSync with a VPC endpoint


B. AWS Direct Connect

C. AWS Snowball Edge Storage Optimized **Most Voted**

D. AWS Storage Gateway

Correct Answer: C

Community vote distribution



Comments

Cyberkayu **Highly Voted** 1 year, 5 months ago

Selected Answer: C

90% utilization of the bandwidth = they discouraged the use of internet bandwidth for uploading, go seek for offline data seeding to AWS method
upvoted 8 times

NSA_Poker 1 year ago

data centers don't "discouraged the use of internet bandwidth."
upvoted 1 times

TariqKipkemei **Highly Voted** 1 year, 6 months ago

Selected Answer: C

50 TB of data to AWS within 2 weeks = Snowball Edge Storage Optimized
upvoted 5 times

KennethNg923 **Most Recent** 11 months, 3 weeks ago

Selected Answer: C

50 TB of data to AWS within 2 weeks + 90% utilization already be used -> snowball

upvoted 2 times

NSA_Poker 1 year ago

Selected Answer: B

(B) is correct.

Direct Connect is "A dedicated connection [is] made through a 1-Gbps, 10-Gbps, or 100-Gbps Ethernet port dedicated to a single customer." DX uses 802.1Q VLANs providing a dedicated private network connection to AWS. At 1-Gbps, transfer takes less than 5 days; at 100-Gbps it takes less than 67 minutes. Since it's a data center & not an oil rig in the middle of the Gulf of Mexico, data center should be able to get this service.

(C) is incorrect

"Consider Snowball Edge if you need to run computing in rugged, austere, mobile, or disconnected (or intermittently connected) environments. Also consider it for large-scale data transfers and migrations when bandwidth is not available for use of a high-speed online transfer service, such as AWS DataSync."

upvoted 1 times

JA2018 6 months, 2 weeks ago

hi bro, setting a new DX will usually takes longer than 2 weeks..... and we have not even started any Data transfer over the wire yet.

upvoted 1 times

NSA_Poker 1 year ago

B > A > C > D

upvoted 1 times

xBUGx 1 year, 2 months ago

Assuming vpn is 1Gbps, it can still transfer 50TB with in 5days with only 10% bandwidth available

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: C

A DataSync is for data

B Direct connect takes longer than 2 weeks

D StorageGateway is useless without more contex

C is only remaining choice.

upvoted 3 times

NSA_Poker 1 year ago

(A) The problem is not that DataSync deals with data (we have to 'transfer 50 TB of data'). The issue is with the VPC endpoint; it doesn't increase our bandwidth. DataSync task is capable fully utilizing 10-Gbps over a network link between your on-premises environment and AWS.

(B) works between 67-minutes & 5 days.

(C) Works but is NOT what I would recommend to a data center.

(D) over VPN connection that is 90% utilized, takes more than 6 weeks.

upvoted 1 times

ftaws 1 year, 5 months ago

Not mentioned network bandwidth. How we know that?

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #660

Topic 1

A company hosts an application on Amazon EC2 On-Demand Instances in an Auto Scaling group. Application peak hours occur at the same time each day. Application users report slow application performance at the start of peak hours. The application performs normally 2-3 hours after peak hours begin. The company wants to ensure that the application works properly at the start of peak hours.

Which solution will meet these requirements?

- A. Configure an Application Load Balancer to distribute traffic properly to the instances.
- B. Configure a dynamic scaling policy for the Auto Scaling group to launch new instances based on memory utilization.
- C. Configure a dynamic scaling policy for the Auto Scaling group to launch new instances based on CPU utilization.
- D. Configure a scheduled scaling policy for the Auto Scaling group to launch new instances before peak hours. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

ZZZ_Sleep **Highly Voted** 11 months, 3 weeks ago

Selected Answer: D

occur at the same time each day = predictable

So, scheduled scaling policy, Answer is D.

Dynamic scaling policy work for unpredictable

upvoted 7 times

Arnaud92 **Highly Voted** 1 year ago

D. The application performs normally 2-3 hours after peak hours begin is a key! (schedule policy)

upvoted 5 times

awsgeek75 **Most Recent** 11 months ago

Selected Answer: D

ABC won't solve the performance issues at the start of peak hours

ABC won't solve the performance issues at the start of peak hours.

D ensure that application is ready for use during the peak hours by scheduling an early launch

upvoted 2 times

TariqKipkemei 1 year ago

Selected Answer: D

Technically both dynamic and scheduled scaling would work but there is strict requirement for the application to work properly at the start of peak hours and no mention of cost.

So scheduled scaling policy it is.

upvoted 4 times

TOR_0511 1 year ago

Selected Answer: D

Application users report slow application performance at the start of peak hours. The company wants to ensure that the application works properly at the start of peak hours

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #661

Topic 1

A company runs applications on AWS that connect to the company's Amazon RDS database. The applications scale on weekends and at peak times of the year. The company wants to scale the database more effectively for its applications that connect to the database.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use Amazon DynamoDB with connection pooling with a target group configuration for the database. Change the applications to use the DynamoDB endpoint.

B. Use Amazon RDS Proxy with a target group for the database. Change the applications to use the RDS Proxy endpoint.

Most Voted

C. Use a custom proxy that runs on Amazon EC2 as an intermediary to the database. Change the applications to use the custom proxy endpoint.

D. Use an AWS Lambda function to provide connection pooling with a target group configuration for the database. Change the applications to use the Lambda function.

Correct Answer: B

Community vote distribution

B (100%)

Comments

TariqKipkemei Highly Voted 1 year ago

Selected Answer: B

Amazon RDS Proxy is a fully managed, highly available database proxy for Amazon Relational Database Service (RDS) that makes applications more resilient to database failures. Many applications, including those built on modern serverless architectures, can have a large number of open connections to the database server and may open and close database connections at a high rate, exhausting database memory and compute resources. Amazon RDS Proxy allows applications to pool and share connections established with the database, improving database efficiency and application scalability. With RDS Proxy, failover times for Aurora and RDS databases are reduced by up to 66%

upvoted 7 times

TOR_0511 Highly Voted 1 year ago

Selected Answer: B

A out because DynamoDB is a NoSQL DB

B As the question is referring about DB connections so this option has the LEAST operational overhead
upvoted 5 times

awsgeek75 Most Recent 11 months ago

Selected Answer: B

A: DynamoDB != RDS

C: Total nonsense

D: Lambda for providing connection pooling sound impractical if not impossible. Would be fun to watch someone do this though...

B RDS Proxy is specifically made for connection pooling.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #662

Topic 1

A company uses AWS Cost Explorer to monitor its AWS costs. The company notices that Amazon Elastic Block Store (Amazon EBS) storage and snapshot costs increase every month. However, the company does not purchase additional EBS storage every month. The company wants to optimize monthly costs for its current storage usage.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use logs in Amazon CloudWatch Logs to monitor the storage utilization of Amazon EBS. Use Amazon EBS Elastic Volumes to reduce the size of the EBS volumes.
- B. Use a custom script to monitor space usage. Use Amazon EBS Elastic Volumes to reduce the size of the EBS volumes.
- C. Delete all expired and unused snapshots to reduce snapshot costs.
- D. Delete all nonessential snapshots. Use Amazon Data Lifecycle Manager to create and manage the snapshots according to the company's snapshot policy requirements. **Most Voted**

Correct Answer: D*Community vote distribution*

D (100%)

Comments**t0nx** **Highly Voted** 1 year, 6 months ago**Selected Answer: D**

This option involves managing snapshots efficiently to optimize costs with minimal operational overhead.

Delete all nonessential snapshots: This reduces costs by eliminating unnecessary snapshot storage. Use Amazon Data Lifecycle Manager (DLM): DLM can automate the creation and deletion of snapshots based on defined policies. This reduces operational overhead by automating snapshot management according to the company's snapshot policy requirements.

upvoted 8 times

JA2018 **Most Recent** 6 months, 2 weeks ago**Selected Answer: D**

Allows the user to manage snapshots efficiently for cost optimization with the LEAST operational overhead.

#1: Delete all nonessential snapshots: This reduces costs by eliminating unnecessary storage used by such snapshots.

#2: Amazon Data Lifecycle Manager (DLM): This AWS managed service can automate the creation and deletion process of snapshots based on defined policies. Automating snapshot management according to snapshot policy requirements helps to cut down operational overhead .

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: D

Least operational overhead for your snapshot management is
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/snapshot-lifecycle.html>

C will just do it once but assuming they want an ongoing solution.

A: It will help with EBS size but won't fix the snapshot problems

B: Same as A, nothing to do with snapshots

upvoted 4 times

xBUGx 1 year, 2 months ago

Q says The company wants to optimize monthly costs for its current storage usage. I think they only want to do it once?

upvoted 1 times

KennethNg923 11 months, 3 weeks ago

I think "optimize monthly costs" include future monthly cost as well, so D is better than C

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: D

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/snapshot-lifecycle.html>

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #663

Topic 1

A company is developing a new application on AWS. The application consists of an Amazon Elastic Container Service (Amazon ECS) cluster, an Amazon S3 bucket that contains assets for the application, and an Amazon RDS for MySQL database that contains the dataset for the application. The dataset contains sensitive information. The company wants to ensure that only the ECS cluster can access the data in the RDS for MySQL database and the data in the S3 bucket.

Which solution will meet these requirements?

- A. Create a new AWS Key Management Service (AWS KMS) customer managed key to encrypt both the S3 bucket and the RDS for MySQL database. Ensure that the KMS key policy includes encrypt and decrypt permissions for the ECS task execution role. **Most Voted**
- B. Create an AWS Key Management Service (AWS KMS) AWS managed key to encrypt both the S3 bucket and the RDS for MySQL database. Ensure that the S3 bucket policy specifies the ECS task execution role as a user.
- C. Create an S3 bucket policy that restricts bucket access to the ECS task execution role. Create a VPC endpoint for Amazon RDS for MySQL. Update the RDS for MySQL security group to allow access from only the subnets that the ECS cluster will generate tasks in.
- D. Create a VPC endpoint for Amazon RDS for MySQL. Update the RDS for MySQL security group to allow access from only the subnets that the ECS cluster will generate tasks in. Create a VPC endpoint for Amazon S3. Update the S3 bucket policy to allow access from only the S3 VPC endpoint.

Correct Answer: A*Community vote distribution***Comments****pentium75** **Highly Voted** 1 year, 5 months ago**Selected Answer: A**

We're asked to restrict access to both, RDS and S3, to "the ECS cluster" (not to a subnet or endpoint).

Not B: Does not restrict RDS at all. Wording about S3 is unusual.

Not C: Would work for S3, but would allow RDS access from whole subnet which may contain other resources besides the ECS cluster

Not D: Would allow RDS access from whole subnet which may contain other resources besides the ECS cluster. Would allow S3

access from VPC endpoint which might be accessed by other resources besides the ECS cluster.

upvoted 13 times

GOTJ 3 months, 1 week ago

You are right with your option "D" concerns about letting other resources running on the same subnet that your ECS cluster accessing the database. But, likewise, the "task execution role" mentioned in your selected answer "A" is applied to a specific ECS task, not to the whole cluster (requirement)

upvoted 1 times

deacon967 5 months, 1 week ago

A doesn't restrict access. It only encrypts at rest. D restricts access. Encryption and Access Control are 2 different things

upvoted 2 times

t0nx Highly Voted 1 year, 6 months ago

Selected Answer: D

Option D is the most comprehensive solution as it leverages VPC endpoints for both Amazon RDS and Amazon S3, along with proper network-level controls to restrict access to only the necessary resources from the ECS cluster.

upvoted 9 times

awsgeek75 1 year, 4 months ago

D only secures access to RDS and S3, it does not secure the sensitive data inside the RDS and S3.

upvoted 3 times

marcosnamekata Most Recent 1 month, 1 week ago

Selected Answer: D

D- Shure

upvoted 1 times

diazed 1 month, 1 week ago

Selected Answer: C

I think this question is not written correctly.

upvoted 1 times

tch 2 months, 2 weeks ago

Selected Answer: A

But security groups allow you to control inbound and outbound traffic at the instance level, while network ACLs offer similar capabilities at the VPC subnet level.

upvoted 1 times

tch 2 months, 2 weeks ago

look portion of answer D: Update the RDS for MySQL security group to allow access from only the subnets that the ECS cluster will generate tasks in.

"allow access from only the subnets " .. is this correct for security group?

upvoted 1 times

Dantecito 3 months, 3 weeks ago

Selected Answer: C

I think the word "restricts" is confusing. But I interpret it as allow the ECS task execution role to access the S3 bucket and deny everything else.

upvoted 1 times

dragosky 6 months ago

Selected Answer: D

The dataset contains sensitive information, nothing to be done here, just stating that it contains sensitive info.

upvoted 1 times

IA2018 6 months, 2 weeks ago

ONLY 10 6 months, 2 weeks ago

Selected Answer: C

Why not C?

Key point:

The question specifically asks to ensure that only the ECS cluster can access the data, which means restricting access to the S3 bucket and RDS database solely to the ECS task execution role.

Breakdown of option C:

S3 bucket policy: By creating a policy that only allows the ECS task execution role to access the S3 bucket, you effectively limit access to the data stored there.

VPC endpoint for RDS: A VPC endpoint creates a private connection to the RDS database, preventing any external access and ensuring that only resources within the VPC can connect.

RDS security group update: Further restricting access by configuring the RDS security group to only allow connections from the subnets where the ECS cluster will launch tasks.

upvoted 1 times

JA2018 6 months, 2 weeks ago

Why the other options are not as suitable:

Option A and B:

While using a KMS key to encrypt data is a good practice, it doesn't directly control access to the data. Attaching the KMS key to the ECS task execution role in the policy would still allow any entity with the key to decrypt the data, not just the ECS cluster.

upvoted 1 times

JA2018 6 months, 2 weeks ago

Option D:

This option only addresses the RDS access by creating a VPC endpoint, but it doesn't restrict access to the S3 bucket effectively, which is also mentioned in the question

upvoted 1 times

XXXXXINN 8 months ago

I cannot believe how many people vote A.

the questions is asking only allow ECS cluster access RDS and access to S3.

2 keys here: 1. security group is usually used to security access between RDS and ECS cluster 2. access data in S3 securely, immediately, we should think about S3 VPC Gateway endpoints because this secures the traffic only travel via private network.

Answer A is just talking about encrypt data at rest, and that is not what the question is asking about

upvoted 3 times

MandAsh 11 months, 3 weeks ago

After reading comments changed to A. D will not protect data at rest it will only give n/w level security

upvoted 1 times

bujuman 1 year, 1 month ago

Selected Answer: A

According to me "The dataset contains sensitive information" is the main information that motivate the real requirement which is "The company wants to ensure that only the ECS cluster can access the data in the RDS for MySQL database and the data in the S3 bucket". So we have to take these two assertions into consideration.

And knowing that, as S3 default encryption capabilities, RDS Mysql DB Instance encryption is not active by default (check this link for details <https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Overview.Encryption.html>), option A is the best option to meet the requirements of accessing the datasets and the assets only from ECS cluster tasks and preserve, at the same time, data confidentiality and integrity. In other words, option A is the best one to ensure the data protection at REST for S3 and RDS and only accessed by ECS cluster.

upvoted 3 times

Hung23 1 year, 1 month ago

Selected Answer: C

Selected Answer: C

Try to chat GPT Please
upvoted 1 times

seetpt 1 year, 3 months ago

Selected Answer: A

A seems right
upvoted 2 times

[Removed] 1 year, 4 months ago

Selected Answer: A

Vote for A. Keywords: "sensitive information" and "data in..."
D: only network control, can't control data access on sensitive information.
upvoted 4 times

Marco_St 1 year, 5 months ago

Selected Answer: C

I did not get how does D achieves the only access from ECS cluster to S3 VPC endpoint.
upvoted 1 times

1rob 1 year, 5 months ago

Selected Answer: A

A; When Only the ECS task execution role is able to encrypt and decrypt the data in the RDS and in the S3 bucket by means of the KMS key policy, you ensure that nothing else can read or modify the data.
B: this answer doesn't state that only the ECS cluster can reach the data.
C: Creating a VPC endpoint for RDS does not mean that only the ECS cluster can reach the data
D: The S3 VPC endpoint does not guarantee that only the ECS cluster can reach the data. Also allowing a subnet to have access to the RDS sounds too open to me
upvoted 5 times

Min_93 1 year, 5 months ago

Options A and B involve using AWS Key Management Service (AWS KMS) for encryption but do not directly address the requirement to restrict access to the ECS cluster. Option C is not the most direct approach for restricting access to the RDS database, as it focuses on the S3 bucket.

Therefore, option D is the most appropriate solution for ensuring that only the ECS cluster can access the data in the RDS for MySQL database and the data in the S3 bucket.
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #664

Topic 1

A company has a web application that runs on premises. The application experiences latency issues during peak hours. The latency issues occur twice each month. At the start of a latency issue, the application's CPU utilization immediately increases to 10 times its normal amount.


The company wants to migrate the application to AWS to improve latency. The company also wants to scale the application automatically when application demand increases. The company will use AWS Elastic Beanstalk for application deployment.

Which solution will meet these requirements?

- A. Configure an Elastic Beanstalk environment to use burstable performance instances in unlimited mode. Configure the environment to scale based on requests. **Most Voted**
- B. Configure an Elastic Beanstalk environment to use compute optimized instances. Configure the environment to scale based on requests.
- C. Configure an Elastic Beanstalk environment to use compute optimized instances. Configure the environment to scale on a schedule.
- D. Configure an Elastic Beanstalk environment to use burstable performance instances in unlimited mode. Configure the environment to scale on predictive metrics.

Correct Answer: A

Community vote distribution



Comments

pentium75 **Highly Voted** 1 year, 5 months ago

Selected Answer: A

"Scale on predictive metrics" does not sound like something that Beanstalk can do. In EC2 you can create a "predictive scaling policy", but apparently this is not supported by Beanstalk. That would rule out D.

We have no indication that the application is CPU-intensive in general. If CPU utilization "increases to 10 times its normal amount" then the "normal amount" cannot be higher than 10 %. This would rule out B and C.

upvoted 11 times

LemonCremlin **Highly Voted** 1 year, 6 months ago

LeonGremlin 1 year, 6 months ago

Selected Answer: D

Burstable Performance Instances (T3 or T3a): These instances are designed for burstable workloads and provide a baseline level of CPU performance with the ability to burst above that baseline when needed. Bursting is particularly beneficial for handling sudden spikes in CPU utilization, such as those described in the scenario.

Unlimited Mode: Enabling "unlimited" mode allows instances to burst beyond their baseline performance without accumulating CPU credits. This is important for handling sudden and sustained increases in CPU utilization during peak hours.

Scale on Predictive Metrics: Configuring the environment to scale on predictive metrics allows AWS Elastic Beanstalk to proactively adjust the number of instances based on anticipated demand. This can help ensure that the environment is scaled up before the latency issues occur, addressing them in advance.

upvoted 10 times

ftaws 1 year, 5 months ago

Traffic is "immediately increases". We can't predict and can not use Predictive Metrics.
And requirement need auto scaling

upvoted 2 times

Salilgen **Most Recent** 5 months, 1 week ago

Selected Answer: A

IMO answer is A.

"The Auto Scaling group uses two Amazon CloudWatch alarms to trigger scaling operations. The default triggers scale when the average outbound network traffic from each instance is higher than 6 MiB or lower than 2 MiB over a period of five minutes. To use Auto Scaling effectively, configure triggers that are appropriate for your application, instance type, and service requirements. You can scale based on several statistics including latency, disk I/O, CPU utilization, and request count."
This is dynamic scaling.

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/burstable-performance-instances.html>

<https://docs.aws.amazon.com/elasticbeanstalk/latest/dg/using-features-managing-env-types.html>

<https://docs.aws.amazon.com/elasticbeanstalk/latest/dg/using-features-managing-as.html>

upvoted 1 times

3c6417b 12 months ago

Selected Answer: B

Explain to me why it's not B?

upvoted 1 times

LeonSauveterre 5 months, 3 weeks ago

B is incorrect because CPU usage percentage rising does not necessarily have anything to do with computing methods. Especially during peak hours, the jobs are coming in hot, but it's quite possible that each job can be done within seconds. It's just too many jobs, and that's not something compute optimized instances can tackle.

upvoted 1 times

Gape4 11 months, 1 week ago

I have the same question.

upvoted 1 times

sandordini 1 year, 1 month ago

Selected Answer: A

D - No such service as Elastic Beanstalk Predictive Scaling, And even if there was, no historical data in AWS for an application we are just about to migrate to AWS. Therefore: A

upvoted 6 times

lenotc 1 year, 2 months ago

Selected Answer: A

D is incorrect Predictive scaling not fit

upvoted 3 times

awsgeek75 1 year, 4 months ago

For those voting D, predictive scaling analyses historic data to predict the scaling needs. This scenario is a migration scenario so

there won't be any historic data which is why D won't work. A (burst) is the only option after migration.

upvoted 5 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

BC are compute optimised instances which don't solve 10x CPU issues at start of the latency.

AD are burstable performance which will help with 10x increase CPU usage

D is not an available feature of Elastic Beanstalk (yet) or I cannot find it in config/docs. Happy to be corrected

A makes sense due to burst performance. Scale based on requests is possible and I'm assuming that latency is related to requests.

upvoted 6 times

1rob 1 year, 5 months ago

Selected Answer: A

Following <https://docs.aws.amazon.com/elasticbeanstalk/latest/dg/using-features.managing.as.html> I see: " You can scale based on several statistics including latency, disk I/O, CPU utilization, and request count. " So no 'scale on predictive metrics, so D is not okay.

Also, the company also wants to scale the application automatically when application demand increases, so scale on a schedule is not appropriate here. So C is not okay.

Burstable performance instances in unlimited mode can sustain high CPU utilization for any period of time whenever required, so an immediate demand of CPU resources is 'covered'. So I go for A.

upvoted 4 times

Min_93 1 year, 5 months ago

Selected Answer: D

Option A, which suggests using burstable performance instances in unlimited mode, is appropriate. However, option D is more specific to the requirement of scaling based on predictive metrics, which is crucial for handling the latency issues that occur at specific times each month.

Options B and C suggest using compute optimized instances and scaling based on requests or on a schedule. While these options might work for general scalability, they may not address the immediate and intense spikes in CPU utilization that are mentioned in the scenario.

Therefore, option D is the most appropriate solution for improving latency and automatically scaling the application based on predictive metrics using AWS Elastic Beanstalk.

upvoted 4 times

evelynsun 1 year, 5 months ago

Selected Answer: A

This solution meets the requirements because it allows the company to automatically scale the application's CPU capacity based on the number of requests it receives. The burstable performance instances provide high CPU performance when needed, which can help to reduce latency during peak hours.

not D: this solution has some drawbacks. First, it can be expensive to use burstable performance instances in unlimited mode, as the instances are charged per hour. Second, it can be difficult to predict the exact CPU requirements of the application, which can lead to over- or under-provisioning of CPU resources.

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

The company also wants to scale the application automatically when application demand increases = Scale based on requests

upvoted 3 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

B

Question is asking scale based on demand so better scale based on requests. Predictive metrics not defined and may be interpreted differently by many users.

upvoted 2 times

reika1914 1 year, 6 months ago

Selected Answer: D

Given the scenario described, the best solution among the provided options to meet the requirements of migrating the application to AWS, improving latency, and scaling the application automatically during increased demand would be:

D. Configure an Elastic Beanstalk environment to use burstable performance instances in unlimited mode. Configure the environment to scale on predictive metrics.

upvoted 2 times

t0nx 1 year, 6 months ago

Selected Answer: D

In this scenario, the application experiences latency issues during peak hours with a sudden increase in CPU utilization. Using burstable performance instances in unlimited mode allows the application to burst beyond the baseline performance when needed. Configuring the environment to scale on predictive metrics enables proactive scaling based on anticipated increases in demand.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #665

Topic 1

A company has customers located across the world. The company wants to use automation to secure its systems and network infrastructure. The company's security team must be able to track and audit all incremental changes to the infrastructure.

Which solution will meet these requirements?

- A. Use AWS Organizations to set up the infrastructure. Use AWS Config to track changes.
- B. Use AWS CloudFormation to set up the infrastructure. Use AWS Config to track changes. **Most Voted**
- C. Use AWS Organizations to set up the infrastructure. Use AWS Service Catalog to track changes.
- D. Use AWS CloudFormation to set up the infrastructure. Use AWS Service Catalog to track changes.

Correct Answer: B

Community vote distribution

B (100%)

Comments

TariqKipkemei **Highly Voted** 1 year, 6 months ago

Selected Answer: B

use automation to secure its systems and network infrastructure = AWS CloudFormation
track and audit all incremental changes to the infrastructure = AWS Config
upvoted 10 times

Min_93 **Highly Voted** 1 year, 5 months ago

Selected Answer: B

Option B is the most suitable because it combines the benefits of infrastructure as code (CloudFormation) with tracking and auditing capabilities (AWS Config). With CloudFormation, the company can define and deploy its infrastructure in a repeatable and automated way, ensuring consistency and adherence to security standards. AWS Config then complements this by providing visibility into changes and configuration details.
upvoted 5 times

Gape4 **Most Recent** 11 months, 1 week ago

Selected Answer: B

I will go for B
upvoted 2 times

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

Organisations is not really related to this

AWS Service Catalog is like a IaC source control so D is a close option. However B looks more logical.

upvoted 3 times

awsgeek75 1 year, 4 months ago

The difference is in wording: "The company's security team must be able to track and audit all incremental changes to the infrastructure"

If this has to be done BEFORE the deployment then D is the option

If this is AFTER the deployment then B is the option

Hopefully exam will have better language. Good luck!

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #666

Topic 1

A startup company is hosting a website for its customers on an Amazon EC2 instance. The website consists of a stateless Python application and a MySQL database. The website serves only a small amount of traffic. The company is concerned about the reliability of the instance and needs to migrate to a highly available architecture. The company cannot modify the application code.

Which combination of actions should a solutions architect take to achieve high availability for the website? (Choose two.)

A. Provision an internet gateway in each Availability Zone in use.

B. Migrate the database to an Amazon RDS for MySQL Multi-AZ DB instance. **Most Voted**

C. Migrate the database to Amazon DynamoDB, and enable DynamoDB auto scaling.

D. Use AWS DataSync to synchronize the database data across multiple EC2 instances.

E. Create an Application Load Balancer to distribute traffic to an Auto Scaling group of EC2 instances that are distributed across two Availability Zones. **Most Voted**

Correct Answer: BE

Community vote distribution

BE (100%)

Comments

TariqKipkemei **Highly Voted** 12 months ago

Selected Answer: BE

To achieve high availability for the website, Migrate the database to an Amazon RDS for MySQL Multi-AZ DB instance and Create an Application Load Balancer to distribute traffic to an Auto Scaling group of EC2 instances that are distributed across two Availability Zones.

upvoted 9 times

Sergiuss95 **Highly Voted** 7 months, 1 week ago

Selected Answer: BE

I sold my soul to the devil to pass the exam

upvoted 6 times

awsgeek75 **Most Recent** 11 months ago

Selected Answer: BE

B: RDS HA

E: Application HA

C: Company cannot change code so this won't work

A: Does not make sense with other options

D: Makes no sense with other options

upvoted 3 times

Cyberkayu 11 months, 3 weeks ago

A. no failed over mechanism

C. DynamoDB is no SQL, cannot use with MySQL

D. Not HA, just sync/replication tools.

Answer BE.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #667

Topic 1

A company is moving its data and applications to AWS during a multiyear migration project. The company wants to securely access data on Amazon S3 from the company's AWS Region and from the company's on-premises location. The data must not traverse the internet. The company has established an AWS Direct Connect connection between its Region and its on-premises location.

Which solution will meet these requirements?

- A. Create gateway endpoints for Amazon S3. Use the gateway endpoints to securely access the data from the Region and the on-premises location.
- B. Create a gateway in AWS Transit Gateway to access Amazon S3 securely from the Region and the on-premises location.
- C. Create interface endpoints for Amazon S3. Use the interface endpoints to securely access the data from the Region and the on-premises location. **Most Voted**
- D. Use an AWS Key Management Service (AWS KMS) key to access the data securely from the Region and the on-premises location.

Correct Answer: C*Community vote distribution***Comments****Ernestokoro** **Highly Voted** 1 year, 5 months ago

Ans is C: >>You can access Amazon S3 from your VPC using gateway VPC endpoints. After you create the gateway endpoint, you can add it as a target in your route table for traffic destined from your VPC to Amazon S3.

There is no additional charge for using gateway endpoints. Amazon S3 supports both gateway endpoints and interface endpoints. With a gateway endpoint, you can access Amazon S3 from your VPC, without requiring an internet gateway or NAT device for your VPC, and with no additional cost. However, gateway endpoints do not allow access from on-premises networks, from peered VPCs in other AWS Regions, or through a transit gateway. For those scenarios, you must use an interface endpoint, which is available for an additional cost. For more information, see Types of VPC endpoints for Amazon S3 in the Amazon S3 User Guide.

<https://docs.aws.amazon.com/vpc/latest/privatelink/vpc-endpoints-s3.html>

upvoted 11 times

MatAlves **Highly Voted** 8 months, 3 weeks ago

Gateway Endpoint -> only within same VPC

Interface Endpoint -> On-premises (VPN or Direct Connect), or different Region over VPC peering.

upvoted 7 times

bora4motion Most Recent 2 weeks, 5 days ago

Selected Answer: C

to access it from on prem you will need an interface endpoint

upvoted 1 times

chest_jd 1 month, 3 weeks ago

Selected Answer: A

VPC has access to S3 through gateway endpoint or interface endpoint. Gateway is better. If we have Direct connect to VPC from on-premises network we get access to S3.

upvoted 1 times

Gape4 11 months, 1 week ago

Selected Answer: C

Please C

upvoted 2 times

1Alpha1 1 year, 4 months ago

Selected Answer: C

Gateway endpoints do not allow access from on-premises networks, from peered VPCs in other AWS Regions, or through a transit gateway. For those scenarios, you must use an interface endpoint, which is available for an additional cost.

<https://docs.aws.amazon.com/vpc/latest/privatelink/vpc-endpoints-s3.html>

upvoted 3 times

awsgeek75 1 year, 4 months ago

Selected Answer: C

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/privatelink-interface-endpoints.html>

With AWS PrivateLink for Amazon S3, you can provision interface VPC endpoints (interface endpoints) in your virtual private cloud (VPC). These endpoints are directly accessible from applications that are on premises over VPN and AWS Direct Connect, or in a different AWS Region over VPC peering.

upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: C

Not A, Gateway endpoint can be accessed only from inside the VPC it's in

Not B, Transit Gateway alone won't help

Not D, KMS has nothing to do with this

upvoted 4 times

fea9bdf 1 year, 5 months ago

Answer seems to be C

gateway endpoints do not allow access from on-premises networks, from peered VPCs in other AWS Regions, or through a transit gateway. For those scenarios, you must use an interface endpoint, which is available for an additional cost. For more information, see Types of VPC endpoints for Amazon S3 in the Amazon S3 User Guide.

upvoted 4 times

ale_brd_111 1 year, 5 months ago

Selected Answer: C

gateway endpoint uses public ip address even if traffic does not directly route thru the internet, also they are no meant to be used from on-premises. Answer is C

upvoted 3 times

Min_93 1 year, 5 months ago

Selected Answer: C

Options A, B, and D are not the most suitable for the following reasons:

A. Create gateway endpoints for Amazon S3:

Gateway endpoints are used for accessing S3 from within a VPC, but they do not extend connectivity to on-premises locations.

B. Create a gateway in AWS Transit Gateway:

AWS Transit Gateway is designed for routing traffic between VPCs and on-premises networks but is not used as a direct gateway for S3 access.

D. Use an AWS Key Management Service (AWS KMS) key:

AWS KMS is a key management service and does not provide direct access to S3. It's used for managing encryption keys. Therefore, option C, creating interface endpoints for Amazon S3, is the most appropriate solution for securely accessing S3 from both the AWS Region and the on-premises location.

upvoted 2 times

Min_93 1 year, 5 months ago

Gateway endpoints for Amazon S3

Interface endpoints for Amazon S3

In both cases, your network traffic remains on the AWS network.

Use Amazon S3 public IP addresses

Use private IP addresses from your VPC to access Amazon S3

Use the same Amazon S3 DNS names

Require endpoint-specific Amazon S3 DNS names

Do not allow access from on premises

Allow access from on premises

Do not allow access from another AWS Region

Allow access from a VPC in another AWS Region by using VPC peering or AWS Transit Gateway

Not billed

Billed

upvoted 2 times

ftaws 1 year, 5 months ago

Selected Answer: B

Transit Gateway support inter region.

interface gateway not use in S3

upvoted 1 times

Min_93 1 year, 5 months ago

com.amazonaws.ap-southeast-1.s3 amazon Interface

Interface is now available for S3

upvoted 1 times

Beshowasfy 1 year, 5 months ago

Selected Answer: A

GW Endpoint is only for S3 and DynamoDB, interface endpoint for other services so C is wrong

upvoted 3 times

ale_brd_111 1 year, 5 months ago

you can't access gateway endpoint from on-premises

upvoted 3 times

XXXXXINN 8 months ago

but you can via direct connection
upvoted 2 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: C

S3 gateway endpoints do not currently support access from resources in a different Region, different VPC, or from an on-premises (non-AWS) environment.

<https://aws.amazon.com/blogs/architecture/choosing-your-vpc-endpoint-strategy-for-amazon-s3/#:~:text=associated.%20S3%20gateway-,endpoints,-do%20not%20currently>

upvoted 2 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

C

. S3 gateway endpoints do not currently support access from resources in a different Region, different VPC, or from an on-premises (non-AWS) environment. However, if you're willing to manage a complex custom architecture, you can use proxies. In all those scenarios, where access is from resources external to VPC, S3 interface endpoints access S3 in a secure way.

<https://aws.amazon.com/blogs/architecture/choosing-your-vpc-endpoint-strategy-for-amazon-s3/>

upvoted 4 times

VladanO 1 year, 6 months ago

Selected Answer: A

<https://docs.aws.amazon.com/vpc/latest/privatelink/gateway-endpoints.html>

Gateway VPC endpoints provide reliable connectivity to Amazon S3 and DynamoDB without requiring an internet gateway or a NAT device for your VPC.

There is no additional charge for using gateway endpoints.

upvoted 2 times

pentium75 1 year, 5 months ago

You can't use GW endpoint from on-premises

upvoted 2 times

Prahacker69 5 months, 2 weeks ago

You can with direct connect

upvoted 1 times

t0nx 1 year, 6 months ago

Selected Answer: C

CCCCC

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #668

Topic 1

A company created a new organization in AWS Organizations. The organization has multiple accounts for the company's development teams. The development team members use AWS IAM Identity Center (AWS Single Sign-On) to access the accounts. For each of the company's applications, the development teams must use a predefined application name to tag resources that are created.

A solutions architect needs to design a solution that gives the development team the ability to create resources only if the application name tag has an approved value.

Which solution will meet these requirements?

- A. Create an IAM group that has a conditional Allow policy that requires the application name tag to be specified for resources to be created.
- B. Create a cross-account role that has a Deny policy for any resource that has the application name tag.
- C. Create a resource group in AWS Resource Groups to validate that the tags are applied to all resources in all accounts.
- D. Create a tag policy in Organizations that has a list of allowed application names. **Most Voted**

Correct Answer: D*Community vote distribution*

D (100%)

Comments**awsgeek75** **Highly Voted** 11 months ago**Selected Answer: D**

A: Don't think this is possible.

B: Cross account role with deny policy? Never seen anything like this

C: Resource groups have nothing to do with allowed tags

D: Correct https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_tag-policies.html
upvoted 7 times

pentium75 **Most Recent** 11 months, 1 week ago**Selected Answer: D**

Other options don't make sense

upvoted 4 times

m_y_s 12 months ago

Selected Answer: D

A tag policy can also specify that noncompliant tagging operations on specified resource types are enforced. In other words, noncompliant tagging requests on specified resource types are prevented from completing.

upvoted 2 times

Beshowasfy 12 months ago

Selected Answer: D

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_tag-policies.html

upvoted 3 times

SHAAHIBHUSHANAWS 1 year ago

D

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_tag-policies.html

upvoted 2 times

rcptryk 1 year ago

Selected Answer: D

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_tag-policies.html

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #669

Topic 1

A company runs its databases on Amazon RDS for PostgreSQL. The company wants a secure solution to manage the master user password by rotating the password every 30 days.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon EventBridge to schedule a custom AWS Lambda function to rotate the password every 30 days.
- B. Use the modify-db-instance command in the AWS CLI to change the password.
- C. Integrate AWS Secrets Manager with Amazon RDS for PostgreSQL to automate password rotation. **Most Voted**
- D. Integrate AWS Systems Manager Parameter Store with Amazon RDS for PostgreSQL to automate password rotation.

Correct Answer: C

Community vote distribution

C (100%)

Comments

TariqKipkemei **Highly Voted** 12 months ago

Selected Answer: C

password rotation = AWS Secrets Manager
upvoted 10 times

rcptryk **Highly Voted** 1 year ago

Selected Answer: C

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/rds-secrets-manager.html>
upvoted 6 times

awsgeek75 **Most Recent** 11 months ago

Selected Answer: C

"Least operational overhead"

A: Lambda overhead so not correct

B: CLI = overhead

D: Yes, it can be done but requires more work for integration.

C: This is correct way of doing it.

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/rds-secrets-manager.html#rds-secrets-manager-overview>
upvoted 3 times

pentium75 11 months, 1 week ago

Selected Answer: C

Secrets Manager allows that, least overhead
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #670

Topic 1

A company performs tests on an application that uses an Amazon DynamoDB table. The tests run for 4 hours once a week. The company knows how many read and write operations the application performs to the table each second during the tests. The company does not currently use DynamoDB for any other use case. A solutions architect needs to optimize the costs for the table.

Which solution will meet these requirements?

- A. Choose on-demand mode. Update the read and write capacity units appropriately.
- B. Choose provisioned mode. Update the read and write capacity units appropriately. **Most Voted**
- C. Purchase DynamoDB reserved capacity for a 1-year term.
- D. Purchase DynamoDB reserved capacity for a 3-year term.

Correct Answer: B*Community vote distribution***Comments****OSHOAIB** **Highly Voted** 11 months ago**Selected Answer: A**

On-demand mode Option A: On-demand mode is suitable for workloads that are unpredictable or that do not have significant or consistent database traffic. It automatically scales to accommodate workload demands and charges for the read and write throughput that the application consumes. For infrequent testing, this could be cost-effective because you only pay for what you use during the testing period and don't incur costs when the table is not being accessed.

Whereas for the Option B, if you only run tests once a week for 4 hours, you might pay for unused capacity for the rest of the week unless you manually scale down the capacity after tests are completed, which adds operational overhead.

upvoted 7 times

1Alpha1 **Highly Voted** 10 months ago**Selected Answer: B**

With provisioned capacity mode, you specify the number of reads and writes per second that you expect your application to require, and you are billed based on that. Furthermore if you can forecast your capacity requirements you can also reserve a portion of DynamoDB provisioned capacity and optimize your costs even further.

<https://docs.aws.amazon.com/wellarchitected/latest/serverless-applications-lens/capacity.html>

upvoted 6 times

Dz1990 Most Recent 2 weeks, 1 day ago

Selected Answer: B

The exam writers are being EXTRA sneaky here! While on-demand seems logical for 4 hours/week usage, the key trap is this phrase: "The company knows how many read and write operations the application performs to the table each second." This is AWS exam code for "predictable workload - use provisioned mode."

Known/Predictable Workload + Cost Optimization = Provisioned Mode

Unknown/Variable Workload + Cost Optimization = On-Demand Mode

When they give you specific performance data, they expect you to USE it (provisioned), not ignore it (on-demand). AWS exams reward recognizing keyword signals rather than doing obvious cost calculations.

Answer B because the exam tests if you catch the "predictable workload pattern" over pure math. Classic AWS exam trickery!

upvoted 1 times

yangbo 2 months, 3 weeks ago

Selected Answer: A

if after 4hours change the RCU/WCU ,select B

if after 4hours do not change the RCU/WCU ,select A

upvoted 1 times

mestule 10 months, 1 week ago

Selected Answer: B

DynamoDB On-Demand pricing is about 6.94x the cost of provisioned capacity. If your applications have predictable traffic patterns and you don't mind spending the time to understand those patterns, using DynamoDB's provisioned throughput capacity can save you money.

Also can't set any capacity units for on-demand mode, so A is false in it's premise.

<https://www.serverless.com/blog/dynamodb-on-demand-serverless>

upvoted 5 times

anikolov 10 months, 3 weeks ago

Selected Answer: A

A: is most cost effective (which is a question/requirement) - 4h per week for Tests purpose

upvoted 2 times

awsgeek75 11 months ago

Selected Answer: B

CD are expensive as reserved capacity even with discounts would spend most time in idle mode (over paid, under utilized)

A: On demand is good if you have unpredictable usage,

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ReadWriteCapacityMode.html#HowItWorks.OnDemand>

B: Provisioned is good if you the usage:

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/ProvisionedThroughput.html>

"The company knows how many read and write operations the application performs to the table each second during the tests."

so ideally they can set this as max

upvoted 4 times

theonlyhero 11 months ago

I initially thought it would be A, but when they mentioned "Update the read and write capacity units appropriately." which are automatically set in "on-demand" switched to B

upvoted 2 times

skynetjay 11 months ago

Selected Answer: B

Provisioned Mode should be the answer seeing that the workloads are predictable and DynamoDB isn't used for any other thing.

upvoted 2 times

pentium75 11 months, 1 week ago

Selected Answer: B

Agree with B, on-demand mode might not scale fast enough after the DB has been idle for 164 hours. As they know exactly the number of reads and writes per second, should use provisioned mode, and set capacity to 1 RCU and 1 WCU while the DB is not in use.

upvoted 3 times

meenkaza 11 months, 2 weeks ago

Selected Answer: B

Provisioned Mode (Option B): Provisioned mode allows you to specify the desired read and write capacity units. Since the workload occurs once a week for 4 hours, you can provision the read and write capacity units accordingly to handle the expected load during that time. This can be a more cost-effective option than on-demand pricing for predictable workloads.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #671

Topic 1

A company runs its applications on Amazon EC2 instances. The company performs periodic financial assessments of its AWS costs. The company recently identified unusual spending.

The company needs a solution to prevent unusual spending. The solution must monitor costs and notify responsible stakeholders in the event of unusual spending.

Which solution will meet these requirements?

- A. Use an AWS Budgets template to create a zero spend budget.
- B. Create an AWS Cost Anomaly Detection monitor in the AWS Billing and Cost Management console. **Most Voted**
- C. Create AWS Pricing Calculator estimates for the current running workload pricing details.
- D. Use Amazon CloudWatch to monitor costs and to identify unusual spending.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**meenkaza** **Highly Voted** 11 months, 2 weeks ago**Selected Answer: B**

AWS Cost Anomaly Detection (Option B): AWS Cost Anomaly Detection is designed to automatically detect unusual spending patterns based on machine learning algorithms. It can identify anomalies and send notifications when it detects unexpected changes in spending. This aligns well with the requirement to prevent unusual spending and notify stakeholders.

upvoted 10 times

awsgeek75 **Most Recent** 11 months ago**Selected Answer: B**

Unusual spending = Cost anomaly hence B

upvoted 4 times

pentium75 11 months, 1 week ago**Selected Answer: B**

<https://aws.amazon.com/aws-cost-management/aws-cost-anomaly-detection/>
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #672

Topic 1


A marketing company receives a large amount of new clickstream data in Amazon S3 from a marketing campaign. The company needs to analyze the clickstream data in Amazon S3 quickly. Then the company needs to determine whether to process the data further in the data pipeline.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create external tables in a Spark catalog. Configure jobs in AWS Glue to query the data.
- B. Configure an AWS Glue crawler to crawl the data. Configure Amazon Athena to query the data. **Most Voted**
- C. Create external tables in a Hive metastore. Configure Spark jobs in Amazon EMR to query the data.
- D. Configure an AWS Glue crawler to crawl the data. Configure Amazon Kinesis Data Analytics to use SQL to query the data.

Correct Answer: B

Community vote distribution



Comments

meenkaza **Highly Voted** 1 year, 5 months ago

Selected Answer: B

AWS Glue with Athena (Option B): AWS Glue is a fully managed extract, transform, and load (ETL) service, and Athena is a serverless query service that allows you to analyze data directly in Amazon S3 using SQL queries. By configuring an AWS Glue crawler to crawl the data, you can create a schema for the data, and then use Athena to query the data directly without the need to load it into a separate database. This minimizes operational overhead.

upvoted 7 times

pentium75 **Highly Voted** 1 year, 5 months ago

Selected Answer: B

Neither Glue nor EMR nor Kinesis are used "to query the data"

upvoted 5 times

Johnoppong101 **Most Recent** 9 months, 3 weeks ago

Selected Answer: B

You've come a loooong way...keep going...

Kinesis Data Analytics applications continuously read and process streaming data in real time.

Kinesis Data Analytics applications continuously read and process streaming data in real time. Data is already at rest in S3. So Athena.

<https://docs.aws.amazon.com/kinesisanalytics/latest/dev/how-it-works.html>
upvoted 2 times

bogdannb 11 months, 1 week ago

Selected Answer: D

It says to quickly analyze the data, Athena can't do it so it's D
upvoted 2 times

JA2018 6 months, 2 weeks ago

But Athena can directly access S3
upvoted 1 times

JA2018 6 months, 2 weeks ago

requirement: LEAST operational overhead
upvoted 1 times

OSHOAIB 1 year, 5 months ago

Selected Answer: B

Option B - leverages serverless services that minimise management tasks and allows the company to focus on querying and analysing the data with the LEAST operational overhead.
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #673

Topic 1


A company runs an SMB file server in its data center. The file server stores large files that the company frequently accesses for up to 7 days after the file creation date. After 7 days, the company needs to be able to access the files with a maximum retrieval time of 24 hours.

Which solution will meet these requirements?

- A. Use AWS DataSync to copy data that is older than 7 days from the SMB file server to AWS.
- B. Create an Amazon S3 File Gateway to increase the company's storage space. Create an S3 Lifecycle policy to transition the data to S3 Glacier Deep Archive after 7 days. **Most Voted**
- C. Create an Amazon FSx File Gateway to increase the company's storage space. Create an Amazon S3 Lifecycle policy to transition the data after 7 days.
- D. Configure access to Amazon S3 for each user. Create an S3 Lifecycle policy to transition the data to S3 Glacier Flexible Retrieval after 7 days.

Correct Answer: B

Community vote distribution



Comments

awsgeek75 **Highly Voted** 11 months ago

Selected Answer: B

A: DataSync is not used for this
 C: FSx File Gateway requires NFS on both sides so won't work with S3
 D: Doesn't say how to transfer data to S3

B: S3 File Gateway will connect SMB to S3. Lifecycle policy will move objects to S3 Glacier Deep Archive which support 12 hours retrieval

<https://aws.amazon.com/blogs/aws/new-amazon-s3-storage-class-glacier-deep-archive/>

upvoted 8 times

PegasusForever **Highly Voted** 11 months, 2 weeks ago

Answer is B, Amazon S3 File Gateway supports SMB and NFS, Amazon FSx File Gateway SMB for windows workloads.

upvoted 6 times

NayeraB Most Recent 9 months, 3 weeks ago

It feels like C is there just to mess with everyone
upvoted 5 times

pentium75 11 months, 1 week ago

Selected Answer: B

Not C because FSx File Gateway saves files in FSx for Windows file server, not S3.
Not D because users should access the files via SMB
upvoted 4 times

chickenmf 9 months, 1 week ago

"FSx File Gateway saves files in FSx for Windows File Server, not S3"
-- me spreading misinformation on the Internet >:)
upvoted 1 times

chickenmf 9 months, 1 week ago

While it is optimized for compatibility with Windows environments, the files stored in Amazon S3 through the FSx File Gateway are not limited to Windows-only access.
upvoted 2 times

cciesam 11 months, 1 week ago

Selected Answer: B

S3 file gateway supports SMB and S3 Glacier Deep Archive can retrieve data within 12 hours.
<https://aws.amazon.com/storagegateway/file/s3/>

<https://docs.aws.amazon.com/prescriptive-guidance/latest/backup-recovery/amazon-s3-glacier.html>
upvoted 5 times

Roger_Liu 11 months, 1 week ago

Selected Answer: B

I prefer to choose Amazon S3 File Gateway.
<https://docs.aws.amazon.com/filegateway/latest/files3/file-gateway-concepts.html>
upvoted 5 times

meenkaza 11 months, 2 weeks ago

Selected Answer: C

Amazon FSx File Gateway with S3 Lifecycle policy (Option C): Amazon FSx is a fully managed file storage service, and with a File Gateway, it allows seamless integration between on-premises file servers and AWS storage. By creating an Amazon FSx File Gateway and implementing an S3 Lifecycle policy to transition data to S3 after 7 days, you can achieve the desired storage and retrieval characteristics.
upvoted 3 times

pentium75 11 months, 1 week ago

Wrong. An "FSx File Gateway" stores the files on AWS side in FSx for Windows file server, NOT in S3. Thus you can't apply the "S3 Lifecycle Policy".
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #674

Topic 1

A company runs a web application on Amazon EC2 instances in an Auto Scaling group. The application uses a database that runs on an Amazon RDS for PostgreSQL DB instance. The application performs slowly when traffic increases. The database experiences a heavy read load during periods of high traffic.

Which actions should a solutions architect take to resolve these performance issues? (Choose two.)

A. Turn on auto scaling for the DB instance.

B. Create a read replica for the DB instance. Configure the application to send read traffic to the read replica. **Most Voted**


C. Convert the DB instance to a Multi-AZ DB instance deployment. Configure the application to send read traffic to the standby DB instance.

D. Create an Amazon ElastiCache cluster. Configure the application to cache query results in the ElastiCache cluster. **Most Voted**

E. Configure the Auto Scaling group subnets to ensure that the EC2 instances are provisioned in the same Availability Zone as the DB instance.

Correct Answer: BD

Community vote distribution



Comments

awsgeek75 **Highly Voted** 11 months ago

Selected Answer: BD

A: RDS DB instance Autoscaling is not a thing
 C: You cannot read from standby even if this was done.
 E: Does not solve any problem

Correct answer
 B: Read replicas distribute load and help improving performance
 D: Caching of any kind will help with performance

Remember: " The database experiences a heavy read load during periods of high traffic."

upvoted 10 times

06042022 **Highly Voted** 11 months ago

00042022 Highly Voted 11 months ago

Selected Answer: BD

By creating a read replica, you offload read traffic from the primary DB instance to the replica, distributing the load and improving overall performance during periods of heavy read traffic.

Amazon ElastiCache can be used to cache frequently accessed data, reducing the load on the database. This is particularly effective for read-heavy workloads, as it allows the application to retrieve data from the cache rather than making repeated database queries.

upvoted 5 times

xBUGx Most Recent 8 months, 2 weeks ago

Selected Answer: BD

RDS auto scaling helps capacity issue, not heavy read workload issue.

upvoted 3 times

Tekk97 11 months, 1 week ago

i think we need Multi az DB, with ElastiCache

upvoted 1 times

pentium75 11 months, 1 week ago

Selected Answer: BD

Not A - There is no such thing as "auto scaling for a DB instance". There is automatic storage scaling, but storage is not the issue here.

B - Yes, read replica will help with "heavy read load"

Not C - "send read traffic to the standby DB instance" does not work

D - "Configure the application ..." might be a bit simplified, but ElastiCache helps with read load

Not E - That might have impact on latency, but not on database load; and all instances in same AZ would be against WAF

upvoted 5 times

OSHOAIB 11 months ago

Amazon RDS does support Storage Auto Scaling :)

<https://aws.amazon.com/about-aws/whats-new/2019/06/rds-storage-auto-scaling/>

upvoted 1 times

awsgeek75 11 months ago

Storage auto scaling is not same as instance autoscaling. Storage is not a problem here.

upvoted 3 times

Riajul 11 months, 2 weeks ago

Selected Answer: AB

A and B should be most correct ans

upvoted 3 times

awsgeek75 10 months, 3 weeks ago

A is autoscaling for DB, it won't fix read problem.

upvoted 2 times

Riajul 11 months, 2 weeks ago

Should be A and B

upvoted 1 times

meenkaza 11 months, 2 weeks ago

Selected Answer: BD

B. Create a read replica for the DB instance. Configure the application to send read traffic to the read replica.

By creating a read replica, you offload read traffic from the primary DB instance to the replica, distributing the load and improving overall performance during periods of heavy read traffic.

D. Create an Amazon ElastiCache cluster. Configure the application to cache query results in the ElastiCache cluster.

Amazon ElastiCache can be used to cache frequently accessed data, reducing the load on the database. This is particularly effective for read-heavy workloads, as it allows the application to retrieve data from the cache rather than making repeated database queries.

upvoted 5 times

pentium75 11 months, 1 week ago

ElastiCache requires application changes, "the solutions architect" cannot simply "configure the application to cache query results".

upvoted 3 times

pentium75 11 months, 1 week ago

On second thought, this might still be correct.

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/creating-elasticache-cluster-with-RDS-settings.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #675

Topic 1


A company uses Amazon EC2 instances and Amazon Elastic Block Store (Amazon EBS) volumes to run an application. The company creates one snapshot of each EBS volume every day to meet compliance requirements. The company wants to implement an architecture that prevents the accidental deletion of EBS volume snapshots. The solution must not change the administrative rights of the storage administrator user.

Which solution will meet these requirements with the LEAST administrative effort?

- A. Create an IAM role that has permission to delete snapshots. Attach the role to a new EC2 instance. Use the AWS CLI from the new EC2 instance to delete snapshots.
- B. Create an IAM policy that denies snapshot deletion. Attach the policy to the storage administrator user.
- C. Add tags to the snapshots. Create retention rules in Recycle Bin for EBS snapshots that have the tags.
- D. Lock the EBS snapshots to prevent deletion. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

meenkaza **Highly Voted** 1 year, 5 months ago

Selected Answer: D

Locking EBS Snapshots (Option D): The "lock" feature in AWS allows you to prevent accidental deletion of resources, including EBS snapshots. This can be set at the snapshot level, providing a straightforward and effective way to meet the requirements without changing the administrative rights of the storage administrator user.

upvoted 9 times

ChhatwaniB **Most Recent** 1 month, 1 week ago

Selected Answer: D

<https://docs.aws.amazon.com/ebs/latest/userguide/ebs-snapshot-lock.html>

upvoted 1 times

maoer78 3 months, 2 weeks ago

Selected Answer: C

Recycle Bin for EBS Snapshots

upvoted 2 times

zdi561 4 months ago

Selected Answer: C

D is not right 'Enable the Snapshot Lock feature on important snapshots to prevent any deletion attempts even the administrator. To prevent accidental deletion of EBS volume snapshots, the most effective method is to enable the "Recycle Bin" feature in AWS, which temporarily stores deleted snapshots for a specified retention period, allowing you to recover them if needed

upvoted 2 times

aquarian_ngc 11 months ago

correct option D

upvoted 1 times

Gape4 11 months, 1 week ago

Selected Answer: D

I will go for D

upvoted 1 times

awsgeek75 1 year, 4 months ago

D: Exactly what a locked EBS snapshot is used for
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-snapshot-lock.html>

upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: D

Typical use case for object lock aka D

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #676

Topic 1

A company's application uses Network Load Balancers, Auto Scaling groups, Amazon EC2 instances, and databases that are deployed in an Amazon VPC. The company wants to capture information about traffic to and from the network interfaces in near real time in its Amazon VPC. The company wants to send the information to Amazon OpenSearch Service for analysis.

Which solution will meet these requirements?

- A. Create a log group in Amazon CloudWatch Logs. Configure VPC Flow Logs to send the log data to the log group. Use Amazon Kinesis Data Streams to stream the logs from the log group to OpenSearch Service.
- B. Create a log group in Amazon CloudWatch Logs. Configure VPC Flow Logs to send the log data to the log group. Use Amazon Kinesis Data Firehose to stream the logs from the log group to OpenSearch Service. **Most Voted**
- C. Create a trail in AWS CloudTrail. Configure VPC Flow Logs to send the log data to the trail. Use Amazon Kinesis Data Streams to stream the logs from the trail to OpenSearch Service.
- D. Create a trail in AWS CloudTrail. Configure VPC Flow Logs to send the log data to the trail. Use Amazon Kinesis Data Firehose to stream the logs from the trail to OpenSearch Service.

Correct Answer: B*Community vote distribution***Comments****pentium75** **Highly Voted** 1 year, 5 months ago**Selected Answer: B**

CloudTrail is for logging administrative actions, we need CloudWatch. We want the data in another AWS service (OpenSearch), not Kinesis, thus we need Firehose, not Streams.

upvoted 9 times

meenkaza **Highly Voted** 1 year, 5 months ago**Selected Answer: B**

Amazon CloudWatch Logs and VPC Flow Logs (Option B): VPC Flow Logs capture information about the IP traffic going to and from network interfaces in a VPC. By configuring VPC Flow Logs to send the log data to a log group in Amazon CloudWatch Logs, you can then use Amazon Kinesis Data Firehose to stream the logs from the log group to Amazon OpenSearch Service for analysis. This approach provides near real-time streaming of logs to the analytics service.

upvoted 5 times

Jacky_S Most Recent 11 months, 2 weeks ago

Selected Answer: A

base on the research, it should be Answer A, because question is asking for a "near real time" which Kinesis Data Stream is offering the data with less than 1 second latency. But Kinese Data Firehost is offering the data with more than 1 second.

<https://docs.aws.amazon.com/opensearch-service/latest/developerguide/integrations.html#integrations-kinesis>

<https://stackoverflow.com/questions/44608274/is-there-any-difference-in-processing-times-between-aws-kinesis-firehose-and-str>

<https://docs.aws.amazon.com/streams/latest/dev/using-other-services-cw-logs.html>

upvoted 1 times

FlyingHawk 5 months, 1 week ago

While Amazon Kinesis Data Streams can stream data, it requires additional setup (e.g., Lambda functions) to process and send logs to OpenSearch Service. This adds complexity compared to Kinesis Data Firehose, which is purpose-built for this use case.

upvoted 2 times

Jacky_S 11 months, 2 weeks ago

base on the research, it should be Answer A, because question is asking for a "near real time" which Kinesis Data Stream is offering the data with less than 1 second latency. But Kinese Data Firehost is offering the data with more than 1 second.

<https://docs.aws.amazon.com/opensearch-service/latest/developerguide/integrations.html#integrations-kinesis>

<https://stackoverflow.com/questions/44608274/is-there-any-difference-in-processing-times-between-aws-kinesis-firehose-and-str>

upvoted 1 times

zinabu 1 year, 2 months ago

log analysis place= aws cloudwatch log

data capturing on the entire vpc=aws flow log

near real time data analysis and send to OpenSearch service= kinesis data fire hose

upvoted 2 times

1Alpha1 1 year, 3 months ago

Selected Answer: B

OpenSearch patterns for CloudWatch Logs:

1) "Near Real Time": CloudWatch logs --> Subscription Filter --> Kinesis Data Firehose --> Amazon OpenSearch (option *B*)

2) "Real Time": CloudWatch logs --> Subscription Filter --> Lambda --> Amazon OpenSearch

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #677


Topic 1

A company is developing an application that will run on a production Amazon Elastic Kubernetes Service (Amazon EKS) cluster. The EKS cluster has managed node groups that are provisioned with On-Demand Instances.

The company needs a dedicated EKS cluster for development work. The company will use the development cluster infrequently to test the resiliency of the application. The EKS cluster must manage all the nodes.

Which solution will meet these requirements MOST cost-effectively?

- A. Create a managed node group that contains only Spot Instances. **Most Voted**
- B. Create two managed node groups. Provision one node group with On-Demand Instances. Provision the second node group with Spot Instances.
- C. Create an Auto Scaling group that has a launch configuration that uses Spot Instances. Configure the user data to add the nodes to the EKS cluster.
- D. Create a managed node group that contains only On-Demand Instances.

Correct Answer: A*Community vote distribution***Comments****pentium75** **Highly Voted** 11 months, 1 week ago**Selected Answer: A**

I think the question is easy to misunderstand, whether you should create the whole setup or just the development cluster. But from the wording ("The [production] EKS cluster has (!) managed node groups ... The company needs a dedicated EKS cluster for development work"), I conclude that we should only create the development cluster.

As this will be used "infrequently" for testing purposes only, and it must be "most cost-effective", I'd go with A - new cluster with "one managed node group that contains only Spot instances".

upvoted 10 times

awsgeek75 11 months ago

The wording of question and options is so confusing. The last line is a throw off also "The EKS cluster must manage all the

nodes" Which EKS cluster? A new one or the existing one.
Both A and B are correct depending on how you decipher the question.
I really hope the exam question uses better language!
upvoted 4 times

Drew3000 8 months, 1 week ago

I hate this question.... I think I will go with B just because wording also. A company is developing an application that "WILL" run on a production Amazon Elastic Kubernetes Service
upvoted 1 times

frmrkc **Highly Voted** 10 months, 1 week ago

Selected Answer: B

This question is convoluted and missing some details.

We need:

- control plane running on on-demand EC2s
- worker nodes running on spot instances

Read this to understand correct solution:

<https://aws.amazon.com/blogs/containers/amazon-eks-now-supports-provisioning-and-managing-ec2-spot-instances-in-managed-node-groups/>

upvoted 5 times

bora4motion **Most Recent** 3 weeks, 6 days ago

Selected Answer: A

this is one stupid question with a lot of fluff. The question is around the dev cluster, the answer is A.

upvoted 1 times

tch 2 months, 2 weeks ago

Selected Answer: B

in the exam.... answer A is too easy...

AWS want the candidate to apply the two workload... the PROD vs DEV...

upvoted 1 times

CTao 3 months, 1 week ago

Selected Answer: A

infrequently + cost effective -> Spot

upvoted 1 times

ieffiong 3 months, 1 week ago

Selected Answer: A

For a development Amazon Elastic Kubernetes Service (Amazon EKS) cluster that is used infrequently to test application resiliency, the most cost-effective solution is to create a managed node group consisting solely of Spot Instances (Option A). Spot Instances offer significant cost savings—up to 90% compared to On-Demand Instances—making them ideal for non-critical, intermittent workloads like development and testing.

upvoted 1 times

Dantecito 3 months, 3 weeks ago

Selected Answer: B

I'll go with B because no matter what production should not go in spot instances

upvoted 1 times

ieffiong 3 months, 1 week ago

This has nothing to do with production instances though. The question is referring to the development instance.

upvoted 1 times

zdi561 4 months ago

Selected Answer: D

Clearly prod and dev are two different envs and two separate EKS cluster. In order to do meaningful resilience test, dev must matches the prod, so D is the answer.

upvoted 1 times

bujuman 7 months, 2 weeks ago

Selected Answer: B

If we look closer to the last requirement "The EKS cluster must manage all the nodes." Option B is the only feasible and cost-effective one.

upvoted 3 times

FlyingHawk 5 months, 1 week ago

I think this requirement means to use managed node group to manage all worker nodes, the nodes for control panel will be managed by AWS, so A is correct

upvoted 1 times

1Alpha1 10 months ago

Selected Answer: A

Based on the document [1], we can know that only self-managed node group can deploy the container on EC2 dedicated hosts. Which means that customer needs to manually create launch template, auto scaling group, and register it to the EKS cluster. The creation process should be same as general EC2 auto scaling creation. For now, EKS managed node group only supports on-demand and spot.

MOST cost-effectively: *Spot Instances*

<https://repost.aws/questions/QUugoX4f1gRHW0MGHRTTHFFFA/how-to-create-eks-cluster-with-dedicated-host-node-group>

upvoted 3 times

anikolov 10 months, 3 weeks ago

Selected Answer: A

"The company will use the development cluster infrequently to test the resiliency of the application" = Spot instances = cost effective

upvoted 2 times

06042022 11 months ago

Selected Answer: B

The keywords are infrequent and resiliency..

This solution allows you to have a mix of On-Demand Instances and Spot Instances within the same EKS cluster. You can use the On-Demand Instances for the development work where you need dedicated resources and then leverage Spot Instances for testing the resiliency of the application. Spot Instances are generally more cost-effective but can be terminated with short notice, so using a combination of On-Demand and Spot Instances provides a balance between cost savings and stability.

Option A (Create a managed node group that contains only Spot Instances) might be cost-effective, but it could introduce potential challenges for tasks that require dedicated resources and might not be the best fit for all scenarios.

upvoted 3 times

mr123dd 11 months ago

Selected Answer: B

The GBT vote A, I know the spot instance is the cheapest, but the question says "dedicated EKS cluster for development", so I vote B

upvoted 2 times

FlyingHawk 5 months, 1 week ago

A dedicated EKS cluster means a separate EKS cluster that is isolated from the production cluster. This ensures that development work does not interfere with the production environment

upvoted 1 times

OSHOAIB 11 months ago

Selected Answer: A

Option A leverages the cost savings of Spot Instances, which is ideal for a development environment where the application is tested infrequently, and there is flexibility in when the nodes can be interrupted. This aligns with the goal of cost-efficiency and takes advantage of EKS's ability to manage the nodes directly.

upvoted 1 times

cciesam 11 months, 1 week ago

Selected Answer: B

B is the best ans.

upvoted 1 times

pentium75 11 months, 1 week ago

Why do you think so?

upvoted 2 times

Naijaboy99 11 months, 2 weeks ago

Option B

upvoted 2 times

pentium75 11 months, 1 week ago

Why do you think so?

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #678

Topic 1

A company stores sensitive data in Amazon S3. A solutions architect needs to create an encryption solution. The company needs to fully control the ability of users to create, rotate, and disable encryption keys with minimal effort for any data that must be encrypted.

Which solution will meet these requirements?

- A. Use default server-side encryption with Amazon S3 managed encryption keys (SSE-S3) to store the sensitive data.
- B. Create a customer managed key by using AWS Key Management Service (AWS KMS). Use the new key to encrypt the S3 objects by using server-side encryption with AWS KMS keys (SSE-KMS). **Most Voted**
- C. Create an AWS managed key by using AWS Key Management Service (AWS KMS). Use the new key to encrypt the S3 objects by using server-side encryption with AWS KMS keys (SSE-KMS).
- D. Download S3 objects to an Amazon EC2 instance. Encrypt the objects by using customer managed keys. Upload the encrypted objects back into Amazon S3.

Correct Answer: B*Community vote distribution*

B (96%)

C

Comments**meenkaza** **Highly Voted** 1 year, 5 months ago**Selected Answer: B**

SSE-KMS with Customer Managed Key (Option B): This option allows you to create a customer managed key using AWS KMS. With a customer managed key, you have full control over key lifecycle management, including the ability to create, rotate, and disable keys with minimal effort. SSE-KMS also integrates with AWS Identity and Access Management (IAM) for fine-grained access control.

upvoted 11 times

MatAlves **Highly Voted** 8 months, 3 weeks ago**Selected Answer: B**

Having both awsgeek75 and pentium75 in the comment section makes me more confident about my own answers.

upvoted 6 times

rubiteb **Most Recent** 1 year, 3 months ago

Selected Answer: C

Customer needs to control the 'user's ability' and not the management of the keys. Option C will prevent users to have this ability.

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

Has to be customer manages to AC are not useful

D is just wrong way of doing this

B give full control to customer while using S3 server side encryption.

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: B

A and C do not allow the company "to fully control the ability of users to create, rotate, and disable encryption keys". D is anything but "minimal effort".

upvoted 4 times

Riajul 1 year, 5 months ago

Selected Answer: B

Option B should be correct

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #679

Topic 1

A company wants to back up its on-premises virtual machines (VMs) to AWS. The company's backup solution exports on-premises backups to an Amazon S3 bucket as objects. The S3 backups must be retained for 30 days and must be automatically deleted after 30 days.

Which combination of steps will meet these requirements? (Choose three.)

- A. Create an S3 bucket that has S3 Object Lock enabled. **Most Voted**
- B. Create an S3 bucket that has object versioning enabled.
- C. Configure a default retention period of 30 days for the objects. **Most Voted**
- D. Configure an S3 Lifecycle policy to protect the objects for 30 days.
- E. Configure an S3 Lifecycle policy to expire the objects after 30 days. **Most Voted**
- F. Configure the backup solution to tag the objects with a 30-day retention period

Correct Answer: ACE*Community vote distribution***Comments****te1973** **Highly Voted** 1 year ago

This is a good example for a completely non-sense AWS exam question. In order to delete the object like requested in the question you need (E). This is required in either versioned or non-versioned buckets. Basically the task is done here. But let's assume we want to make it extra secure and retain the files for 30 days. Then we need object lock (A). You cannot have object lock without versioning (B). You also need to set a retention period then (C). So you either have A,B,C,E or you have E. Choosing exactly 3 options is completely nonsense here. But what do i know.

upvoted 8 times

pentium75 **Highly Voted** 1 year, 5 months ago**Selected Answer:** ACE

In theory, E alone would be enough because the objects are "retained for 30 days" without any configuration as long as no one deletes them. But let's assume that they want us to prevent deletion.

A: Yes, required to prevent deletion. Object Lock requires Versioning, so if we 'create an S3 bucket that has S3 Object Lock enabled' that this also has object versioning enabled. otherwise we would not be able to create it

enabled that this also has object versioning enabled, otherwise we would not be able to create it.

B: No. We need versioning, but we cannot "create" the bucket twice. If we create it "with object lock enabled" then versioning is enabled too, but NOT the other way round (creating it with versioning enabled will not automatically enable object lock).

upvoted 5 times

pentium75 1 year, 5 months ago

C: Yes, "default retention period" specifies how long object lock will be applied to new objects by default, we need this to protect objects from deletion.

D: No, S3 Lifecycle Policy can "transition" or "expire" but not "protect".

E: Yes, this will delete the objects after 30 days (C just removes the object lock after 30 days but does not delete the objects).

F: No, 'tag with a retention period' is not common AWS wording, "tags" are something different in AWS context

upvoted 4 times

Dz1990 **Most Recent** 2 weeks, 1 day ago

Selected Answer: ACE

After analyzing the community debate, I'm sticking with ACE.

AWS Logic: They wouldn't give you two "create bucket" options (A and B) if both were needed. When you "Create bucket with Object Lock enabled" (A), versioning is automatically included - you can't have Object Lock without it.

The sneaky part: AWS is testing if you understand Object Lock includes versioning, not if you'll pick redundant options. If they wanted both A and B, they'd phrase them as separate configuration steps, not both as "create bucket" actions.

Most experienced test-takers chose ACE recognizing this pattern. AWS tests concepts, not redundant configurations.

ACE is correct.

upvoted 1 times

tch 2 months, 3 weeks ago

Selected Answer: ACE

To manage the lifecycle of your objects, create an S3 Lifecycle configuration for your bucket. An S3 Lifecycle configuration is a set of rules that define actions that Amazon S3 applies to a group of objects. There are two types of actions:

- Transition actions – These actions define when objects transition to another storage class. For example, you might choose to transition objects to the S3 Standard-IA storage class 30 days after creating them, or archive objects to the S3 Glacier Flexible Retrieval storage class one year after creating them.

- Expiration actions – These actions define when objects expire. Amazon S3 deletes expired objects on your behalf. For example, you might choose to expire objects after they have been stored for a regulatory compliance period.

upvoted 1 times

tch 2 months, 3 weeks ago

To use S3 Object Lock with a bucket (or objects within a bucket), you must first enable versioning for the bucket, as you won't be able to turn versioning on later. Retention periods and legal holds apply to individual object versions

upvoted 1 times

zdi561 3 months, 4 weeks ago

Selected Answer: ABE

C is not right because a standard S3 object has no default retention period, meaning it can be deleted at any time unless you specifically configure a retention period using the "S3 Object Lock" feature, which allows you to set a custom retention period for individual objects or an entire bucket, with the minimum retention period being one day, B is right because you have to do it, it does not matter you have to do it to lock a object or do it yourself. D is a general description of action A.

upvoted 1 times

omega_coaching 5 months, 3 weeks ago

Selected Answer: ABE

For enabling S3 object lock we need to enable Bucket versioning.. SO and must be ABE

upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: ACE

"Object Lock works only in buckets that have S3 Versioning enabled"

However, we can't have 2 options (A and B) telling to create the bucket. So, A is only possible if versioning is already enabled

However, we can't have 2 options (A and B) telling to create the bucket. So, A is only possible if versioning is already enabled.

We need retention period (C), since this is not a case for legal holds:

"Object Lock provides two ways to manage object retention: retention periods and legal holds."

E - obvious reasons.

Ref. <https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.html>

upvoted 1 times

chwieobjom 9 months, 3 weeks ago

this is shit

upvoted 2 times

mohammadthainat 1 year, 2 months ago

Selected Answer: ACE

1- The S3 backups must be retained for 30 days -->

For that you must enable S3 Object Lock (versioning must be enabled) in Compliance Mode and set Retention Period to 30 days. Thus, to achieve this you need 3 options <A, B and C>

2- The S3 backups must be automatically deleted after 30 days. -->

For that you must Create Lifecycle Rule with action Expire current versions of objects (versioning must be enabled) and set Expiration Period to 30 days. Thus to achieve this you need 2 options <B and E>

 is a must here as both locking the objects and deleting them can't be achieved without it. But, when choosing "A.Create an S3 bucket that has S3 Object Lock enabled." this explicitly indicated that versioning is enabled in your bucket.

upvoted 4 times

awsgeek75 1 year, 4 months ago

Selected Answer: ADE

B: No versioning is required

D: Lifecycle is for transitioning or expiring. There is no protection lifecycle policy

F: No such tag

Enable object lock, retain for 30 days (<https://docs.aws.amazon.com/AmazonS3/latest/userguide/batch-ops-retention-date.html>) and expire after 30 days.

upvoted 5 times

MatAlves 8 months, 3 weeks ago

"Object Lock works only in buckets that have S3 Versioning enabled"

However, I still agree with ACE, since the bucket has already been created, so we can't have 2 answers telling to create the bucket.

And yes, for this case, we need retention period:

"Object Lock provides two ways to manage object retention: retention periods and legal holds."

Ref. <https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.html>

upvoted 3 times

awsgeek75 1 year, 4 months ago

I meant ACE! not ADE!

upvoted 3 times

PegasusForever 1 year, 5 months ago

ABE -> <https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.html>

A. Create an S3 bucket that has S3 Object Lock enabled. -> You set a Retention period of 30 days with this feature.

B. Create an S3 bucket that has object versioning enabled -> Object Lock works only in buckets that have S3 Versioning enabled

E. Configure an S3 Lifecycle policy to expire the objects after 30 days. -> It is valid using the lifecycle policy.

upvoted 2 times

PegasusForever 1 year, 5 months ago

After analyzing the question deeply and reading: <https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.html>, I keep A and B, change E per C.

A. Create an S3 bucket that has S3 Object Lock enabled.

B. Create an S3 bucket that has object versioning enabled.

Change E must be automatically deleted after 30 days(objects will be marked as expired not deleted). per C. Configure a default retention period of 30 days for the objects. It feature delete the object.

upvoted 2 times

PegasusForever 1 year, 5 months ago

Selected Answer: ACE

A. Create an S3 bucket that has S3 Object Lock enabled. Enable the S3 Object Lock feature on S3.

C. Configure a default retention period of 30 days for the objects. To lock the objects for 30 days.

E. Configure an S3 Lifecycle policy to expire the objects after 30 days. -> to delete the objects after 30 days.

upvoted 2 times

cciesam 1 year, 5 months ago

Selected Answer: ACE

ACE is the correct ans.

upvoted 5 times

Riajul 1 year, 5 months ago

Selected Answer: ADE

ADE should be correct

upvoted 1 times

pentium75 1 year, 5 months ago

Why?

S3 Lifecycle Policy can "transition" or "expire" but not "protect"

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-expire-general-considerations.html>

upvoted 2 times

Naijaboy99 1 year, 5 months ago

Correct Answer is A C E

upvoted 2 times

meenkaza 1 year, 5 months ago

Selected Answer: ADE

A. Create an S3 bucket that has S3 Object Lock enabled.

S3 Object Lock provides the ability to enforce retention periods on objects, preventing deletion or modification for a specified duration.

D. Configure an S3 Lifecycle policy to protect the objects for 30 days.

By configuring a lifecycle policy, you can define a transition action to move objects to the S3 Glacier storage class (or any other storage class) after 30 days.

E. Configure an S3 Lifecycle policy to expire the objects after 30 days.

upvoted 1 times

pentium75 1 year, 5 months ago

S3 Lifecycle Policy can "transition" or "expire" but not "protect"

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-expire-general-considerations.html>

upvoted 3 times

tch 2 months, 3 weeks ago

To manage the lifecycle of your objects, create an S3 Lifecycle configuration for your bucket. An S3 Lifecycle configuration is a set of rules that define actions that Amazon S3 applies to a group of objects. There are two types of actions:

- Transition actions – These actions define when objects transition to another storage class. For example, you might choose to transition objects to the S3 Standard-IA storage class 30 days after creating them, or archive objects to the S3 Glacier Flexible Retrieval storage class one year after creating them.
 -
 - Expiration actions – These actions define when objects expire. Amazon S3 deletes expired objects on your behalf. For example, you might choose to expire objects after they have been stored for a regulatory compliance period.
- upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #680

Topic 1

A solutions architect needs to copy files from an Amazon S3 bucket to an Amazon Elastic File System (Amazon EFS) file system and another S3 bucket. The files must be copied continuously. New files are added to the original S3 bucket consistently. The copied files should be overwritten only if the source file changes.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an AWS DataSync location for both the destination S3 bucket and the EFS file system. Create a task for the destination S3 bucket and the EFS file system. Set the transfer mode to transfer only data that has changed. **Most Voted**
- B. Create an AWS Lambda function. Mount the file system to the function. Set up an S3 event notification to invoke the function when files are created and changed in Amazon S3. Configure the function to copy files to the file system and the destination S3 bucket.
- C. Create an AWS DataSync location for both the destination S3 bucket and the EFS file system. Create a task for the destination S3 bucket and the EFS file system. Set the transfer mode to transfer all data.
- D. Launch an Amazon EC2 instance in the same VPC as the file system. Mount the file system. Create a script to routinely synchronize all objects that changed in the origin S3 bucket to the destination S3 bucket and the mounted file system.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**meenkaza** **Highly Voted** 11 months, 2 weeks ago

AWS DataSync (Option A): AWS DataSync is designed for efficient and reliable copying of data between different storage solutions. By setting up an AWS DataSync task with the transfer mode set to transfer only data that has changed, you ensure that only the new or modified files are copied. This minimizes data transfer and operational overhead.

upvoted 5 times

pentium75 11 months, 1 week ago

Actually this is not fully correct:

"By setting up an AWS DataSync task with the transfer mode set to transfer only data that has changed, you ensure that only the new or modified files are copied."

"Transfer only data that has changed ... copies only the data and metadata that differs between the source and destination location."

So, if we have a source with existing items and an empty destination (like in this example), "transfer only data that has changed" will transfer all the existing items though in the true sense of the word they have not "changed".

upvoted 4 times

mohammadthainat Most Recent 8 months, 2 weeks ago

Selected Answer: A

DataSync will do an Initial Scan of both S3 buckets. Identifying Differences. Then, Transferring Changes, so technically DataSync will transfer All the data at first run then it will only transfer newly added/modified objects subsequently.

upvoted 4 times

Kezuko 8 months, 3 weeks ago

Have always did this using B, guess now that I know A is less operational

upvoted 4 times

awsgeek75 11 months ago

Selected Answer: A

BD are more operation overhead although B can work in principle
AC uses managed service to transfer data. A fulfils the requirement of "copied files should be overwritten only if the source file changes" so A is correct. B will just copy regardless of the change

upvoted 2 times

awsgeek75 11 months ago

Meant C will transfer everything and copy data without comparing for change

upvoted 2 times

pentium75 11 months, 1 week ago

Selected Answer: A

Transfer only data that has changed – DataSync copies only the data and metadata that differs between the source and destination location.

Transfer all data – DataSync copies everything in the source to the destination without comparing differences between the locations.

<https://docs.aws.amazon.com/datasync/latest/userguide/configure-metadata.html>

(B would work too but is more "operational overhead.")

upvoted 4 times

cciesam 11 months, 1 week ago

Selected Answer: A

ans: A

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #681

Topic 1


A company uses Amazon EC2 instances and stores data on Amazon Elastic Block Store (Amazon EBS) volumes. The company must ensure that all data is encrypted at rest by using AWS Key Management Service (AWS KMS). The company must be able to control rotation of the encryption keys.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create a customer managed key. Use the key to encrypt the EBS volumes. **Most Voted**
- B. Use an AWS managed key to encrypt the EBS volumes. Use the key to configure automatic key rotation.
- C. Create an external KMS key with imported key material. Use the key to encrypt the EBS volumes.
- D. Use an AWS owned key to encrypt the EBS volumes.

Correct Answer: A

Community vote distribution



Comments

pentium75 **Highly Voted** 11 months, 1 week ago

Selected Answer: A

"Able to control rotation of the encryption keys" = customer managed key (created by AWS but managed by the customer in KMS)

upvoted 5 times

AAbirdy **Highly Voted** 10 months, 4 weeks ago

Selected Answer: A

The company must be able to control rotation of the encryption keys = customer managed key

upvoted 5 times

Chen77 **Most Recent** 3 months, 3 weeks ago

Selected Answer: B

The company must ensure encryption at rest using AWS KMS and must control key rotation with minimal operational overhead.

AWS Managed Keys are automatically managed by AWS, including automatic key rotation every year. This approach eliminates manual key management, reducing administrative effort.

This approach eliminates manual key management, reducing administrative effort.

upvoted 1 times

awsgeek75 11 months ago

Selected Answer: A

"The company must be able to control rotation of the encryption keys."

BD does not allow company owned keys

C is too much operational overhead

upvoted 4 times

dikshya1233 11 months ago

Selected Answer: B

The solution that meets the requirements with the LEAST operational overhead is:

B. Use an AWS managed key to encrypt the EBS volumes. Use the key to configure automatic key rotation.

With AWS managed keys (AWS managed CMKs), AWS takes care of key management tasks, including key rotation. This reduces operational overhead as AWS automatically handles key rotation without requiring manual intervention. It is a convenient option for users who want to ensure encryption at rest with minimal effort in managing encryption keys.

upvoted 2 times

awsgeek75 11 months ago

AWS Managed keys don't meet the requirements "The company must be able to control rotation of the encryption keys."

upvoted 3 times

Shobhit2021 11 months, 1 week ago

Selected Answer: A

A is correct option

upvoted 2 times

fea9bdf 11 months, 1 week ago

Answer is C

Details are on this link below:

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/default-bucket-encryption.html>

Amazon S3 buckets have bucket encryption enabled by default, and new objects are automatically encrypted by using server-side encryption with Amazon S3 managed keys (SSE-S3). This encryption applies to all new objects in your Amazon S3 buckets, and comes at no cost to you.

If you need more control over your encryption keys, such as managing key rotation and access policy grants, you can elect to use server-side encryption with AWS Key Management Service (AWS KMS) keys (SSE-KMS), or dual-layer server-side encryption with AWS KMS keys (DSSE-KMS). For more information about SSE-KMS, see [Specifying server-side encryption with AWS KMS \(SSE-KMS\)](#). For more information about DSSE-KMS, see [Using dual-layer server-side encryption with AWS KMS keys \(DSSE-KMS\)](#).

upvoted 1 times

pentium75 11 months, 1 week ago

How does this relate to answer C? With "imported key material" you cannot "control rotation of the encryption keys" (except by importing new keys). SSE-KMS (as mentioned in your explanation = customer managed key = A

upvoted 2 times

Riajul 11 months, 2 weeks ago

Should be option A

upvoted 2 times

Naijaboy99 11 months, 2 weeks ago

option B is the correct answer with least operational overhead on admins

upvoted 1 times

OSHOAIB 11 months ago

AWS managed keys do allow for automatic rotation, but the company does NOT have control over the rotation - AWS manages this automatically without company intervention

manages this automatically without company intervention.

upvoted 2 times

Naijaboy99 11 months, 2 weeks ago

@meenkaza was right the answer is A

upvoted 3 times

meenkaza 11 months, 2 weeks ago

Selected Answer: A

option A (Create a customer managed key. Use the key to encrypt the EBS volumes) is the most suitable option with the least operational overhead for the given requirements.

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #682

Topic 1


A company needs a solution to enforce data encryption at rest on Amazon EC2 instances. The solution must automatically identify noncompliant resources and enforce compliance policies on findings.

Which solution will meet these requirements with the LEAST administrative overhead?

- A. Use an IAM policy that allows users to create only encrypted Amazon Elastic Block Store (Amazon EBS) volumes. Use AWS Config and AWS Systems Manager to automate the detection and remediation of unencrypted EBS volumes. **Most Voted**
- B. Use AWS Key Management Service (AWS KMS) to manage access to encrypted Amazon Elastic Block Store (Amazon EBS) volumes. Use AWS Lambda and Amazon EventBridge to automate the detection and remediation of unencrypted EBS volumes.
- C. Use Amazon Macie to detect unencrypted Amazon Elastic Block Store (Amazon EBS) volumes. Use AWS Systems Manager Automation rules to automatically encrypt existing and new EBS volumes.
- D. Use Amazon Inspector to detect unencrypted Amazon Elastic Block Store (Amazon EBS) volumes. Use AWS Systems Manager Automation rules to automatically encrypt existing and new EBS volumes.

Correct Answer: A

Community vote distribution



Comments

meenkaza **Highly Voted** 11 months, 2 weeks ago

Selected Answer: A

IAM Policy and AWS Config (Option A): By creating an IAM policy that allows users to create only encrypted EBS volumes, you proactively prevent the creation of unencrypted volumes. Using AWS Config, you can set up rules to detect noncompliant resources, and AWS Systems Manager Automation can be used for automated remediation. This approach provides a proactive and automated solution.

upvoted 13 times

LeonSauveterre **Most Recent** 5 months, 3 weeks ago

Selected Answer: A

Not B - because when you want to detect unencrypted EBS volumes using Lambda and EventBridge, the codes you'll need to compose is so much more than you think. Just leave that whole shebang to AWS Config and then you'll choose A accordingly.

upvoted 1 times

88f8032 7 months, 1 week ago

Selected Answer: B

Isn't B simpler?

upvoted 1 times

awsgeek75 11 months ago

Selected Answer: A

B: Too much work

C: Macie is for PII and sensitive data not for encrypted volumes

D: Inspector for OS patching and vulnerability detections

upvoted 2 times

f2e2419 11 months ago

why not B?

upvoted 1 times

OSHOAIB 11 months ago

Selected Answer: A

Option A - enforces the creation of encrypted volumes via IAM policies and uses AWS Config for detection and AWS Systems Manager for remediation with the LEAST administrative overhead.

upvoted 3 times

pentium75 11 months, 1 week ago

Selected Answer: A

A as exactly described here: <https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/automatically-encrypt-existing-and-new-amazon-ebs-volumes.html>

Not B, that could in theory work but would be massive operational overhead

Not C, Macie detects PII data, not unencrypted volumes

Not D, Inspector detects vulnerabilities, not unencrypted volumes

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #683

Topic 1


A company is migrating its multi-tier on-premises application to AWS. The application consists of a single-node MySQL database and a multi-node web tier. The company must minimize changes to the application during the migration. The company wants to improve application resiliency after the migration.

Which combination of steps will meet these requirements? (Choose two.)

- A. Migrate the web tier to Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer. **Most Voted**
- B. Migrate the database to Amazon EC2 instances in an Auto Scaling group behind a Network Load Balancer.
- C. Migrate the database to an Amazon RDS Multi-AZ deployment. **Most Voted**
- D. Migrate the web tier to an AWS Lambda function.
- E. Migrate the database to an Amazon DynamoDB table.

Correct Answer: AC

Community vote distribution



Comments

meenkaza **Highly Voted** 11 months, 2 weeks ago

Selected Answer: AC

Web Tier Migration (Option A): Migrating the web tier to Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer (ALB) provides horizontal scalability, automatic scaling, and improved resiliency. Auto Scaling helps in managing and maintaining the desired number of EC2 instances based on demand, and the ALB distributes incoming traffic across multiple instances.

Database Migration to Amazon RDS Multi-AZ (Option C): Migrating the database to Amazon RDS in a Multi-AZ deployment provides high availability and automatic failover. In a Multi-AZ deployment, Amazon RDS maintains a standby replica in a different Availability Zone, and in the event of a failure, it automatically promotes the replica to the primary instance. This enhances the resiliency of the database.

upvoted 11 times

pentium75 **Highly Voted** 11 months, 1 week ago

Selected Answer: AC

A - ALB is ideal for web application

B - NLB would work too but ALB is better
C - same functionality as on-premises just with 'improved resiliency'
D - would require significant "changes to the application"
E - would require significant "changes to the application"

upvoted 5 times

fea9bdf Most Recent 11 months, 1 week ago

Also Dynamo DB is noSQL, that can not be an option here

upvoted 3 times

Naijaboy99 11 months, 2 weeks ago

option A C

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #684

Topic 1

A company wants to migrate its web applications from on premises to AWS. The company is located close to the eu-central-1 Region. Because of regulations, the company cannot launch some of its applications in eu-central-1. The company wants to achieve single-digit millisecond latency.

Which solution will meet these requirements?

- A. Deploy the applications in eu-central-1. Extend the company's VPC from eu-central-1 to an edge location in Amazon CloudFront.
- B. Deploy the applications in AWS Local Zones by extending the company's VPC from eu-central-1 to the chosen Local Zone. **Most Voted**
- C. Deploy the applications in eu-central-1. Extend the company's VPC from eu-central-1 to the regional edge caches in Amazon CloudFront.
- D. Deploy the applications in AWS Wavelength Zones by extending the company's VPC from eu-central-1 to the chosen Wavelength Zone.

Correct Answer: B*Community vote distribution***Comments****pentium75** **Highly Voted** 11 months, 1 week ago**Selected Answer: B**

"AWS Local Zones are a type of AWS infrastructure deployment that place compute, storage, database, and other select services closer to large population, industry, and IT centers, enabling you to deliver applications that require single-digit millisecond latency to end-users."

A and C tell us to "deploy the applications in eu-central-1" which is exactly what we're not supposed to do.

AWS Wavelength zones are AWS deployments in CSP's networks, has nothing to do with this question.

https://aws.amazon.com/about-aws/global-infrastructure/localzones/features/?nc1=h_ls

upvoted 9 times

awsgeek75 **Highly Voted** 11 months ago

Selected Answer: B

AC is not right "Because of regulations, the company cannot launch some of its applications in eu-central-1"

D: AWS Wavelength is for mobile network

B: Local Zones can be used to launch apps close to a region but not in a region like EUC1 so this works

upvoted 7 times

bodakrishna **Most Recent** 9 months, 2 weeks ago

Correct B:

AWS Local Zones are an extension of AWS infrastructure and bring AWS services closer to end-users, providing ultra-low latency for applications that require single-digit millisecond latencies. By deploying the applications in AWS Local Zones, the company can meet the latency requirements while also complying with regulations that prevent certain applications from being hosted in the eu-central-1 Region.

upvoted 6 times

OSHOAIB 11 months ago

Selected Answer: B

Option B - AWS Local Zones place AWS compute, storage, database, and other select services closer to end-users. This would allow the company to deploy applications within geographic proximity to eu-central-1 without being directly in the region, potentially meeting regulatory requirements and achieving low latency.

Whereas Option D - AWS Wavelength Zones are designed to provide developers the ability to build applications that deliver single-digit millisecond latencies to MOBILE and connected devices. And it's more focused on 5G Apps and may not be directly relevant to Web Apps hosting.

upvoted 2 times

pdragon1981 11 months, 1 week ago

Selected Answer: B

I would go also for B, was in doubt from B or D but I agree with pentium75 the wavelength zones are not designed for this use case however AWS local zones can provide single-digit milisecond latency as described in the link

<https://aws.amazon.com/about-aws/global-infrastructure/localzones/>

upvoted 2 times

Naijaboy99 11 months, 2 weeks ago

option B

upvoted 4 times

meenkaza 11 months, 2 weeks ago

Selected Answer: D

AWS Wavelength (Option D): AWS Wavelength Zones bring AWS services to the edge of the 5G network, providing ultra-low latency for applications that require single-digit millisecond latencies. Deploying applications in Wavelength Zones allows the company to extend its VPC from the eu-central-1 Region to the chosen Wavelength Zone, providing the required low-latency access.

upvoted 6 times

Roger_Liu 11 months, 1 week ago

It looks like D is correct from diagram in the following url.

<https://docs.aws.amazon.com/wavelength/latest/developecguide/how-wavelengths-work.html>

upvoted 1 times

pentium75 11 months, 1 week ago

"Wavelength Zones are AWS infrastructure deployments that embed AWS compute and storage services within communications service providers' (CSP) 5G networks". They reduce latency for mobile users in the CSP's network, but this is not asked here. Local Zones provide "single-digit millisecond latency".

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #685

Topic 1

A company's ecommerce website has unpredictable traffic and uses AWS Lambda functions to directly access a private Amazon RDS for PostgreSQL DB instance. The company wants to maintain predictable database performance and ensure that the Lambda invocations do not overload the database with too many connections.

What should a solutions architect do to meet these requirements?

- A. Point the client driver at an RDS custom endpoint. Deploy the Lambda functions inside a VPC.
- B. Point the client driver at an RDS proxy endpoint. Deploy the Lambda functions inside a VPC. **Most Voted**
- C. Point the client driver at an RDS custom endpoint. Deploy the Lambda functions outside a VPC.
- D. Point the client driver at an RDS proxy endpoint. Deploy the Lambda functions outside a VPC.

Correct Answer: B

Community vote distribution

B (100%)

Comments

ogerber **Highly Voted** 9 months, 3 weeks ago

Selected Answer: B

Option B.

Reduce number of connection to RDS -> RDS Proxy.

"A Lambda function that's outside of a VPC can't access an RDS instance that's inside a VPC."

<https://repost.aws/knowledge-center/connect-lambda-to-an-rds-instance>

upvoted 10 times

Kezuko **Highly Voted** 8 months, 3 weeks ago

Selected Answer: B

Have to be inside VPC in order to access the RDS instance for Lambda

upvoted 5 times

ogerber **Most Recent** 9 months, 3 weeks ago

Option B.

Reduce number of connection to RDS -> RDS Proxy.

"A Lambda function that's outside of a VPC can't access an RDS instance that's inside a VPC."

<https://repost.aws/knowledge-center/connect-lambda-to-an-rds-instance>

upvoted 4 times

Moon239 10 months ago

Same as question 802 in SAA-C02

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #686

Topic 1

A company is creating an application. The company stores data from tests of the application in multiple on-premises locations.

The company needs to connect the on-premises locations to VPCs in an AWS Region in the AWS Cloud. The number of accounts and VPCs will increase during the next year. The network architecture must simplify the administration of new connections and must provide the ability to scale.

Which solution will meet these requirements with the LEAST administrative overhead?

- A. Create a peering connection between the VPCs. Create a VPN connection between the VPCs and the on-premises locations.
- B. Launch an Amazon EC2 instance. On the instance, include VPN software that uses a VPN connection to connect all VPCs and on-premises locations.
- C. Create a transit gateway. Create VPC attachments for the VPC connections. Create VPN attachments for the on-premises connections. **Most Voted**
- D. Create an AWS Direct Connect connection between the on-premises locations and a central VPC. Connect the central VPC to other VPCs by using peering connections.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**1Alpha1** **Highly Voted** 1 year, 3 months ago**Selected Answer: C**

multiple on-premises locations + increasing number of accounts and VPCs --> connections using *transit gateway*
upvoted 5 times

Andy_09 **Highly Voted** 1 year, 4 months ago

Think C would be the correct answer here.
upvoted 5 times

Tendu **Most Recent** 3 months, 2 weeks ago

Selected Answer: C

Transit Gateway is a way to go in this situation.

upvoted 1 times

LeonSauveterre 5 months, 3 weeks ago

Selected Answer: C

A - OK but for every new VPC or account, you need to manually create new peering connections and VPNs. Too much administrative overhead!

B - Launching EC2s are never an easy job. You not only need to do everything manually mentioned in A, but now you also have to manage the EC2 instance and VPN software. Plus, if the EC2 instance goes down, all connectivity is lost, which is called "single-point of failure".

C - Everything seems cool, it's just uhm lightly higher initial setup cost, but it saves time and money in the long run.

D - Direct Connect provides fast & reliable connections, but just like A & B, as the number of VPCs grows, you'll go crazy due to all those unmanageable nodes.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: C

Anytime I see "the number of VPCs will increase", I immediately look for "transit gateway" as the least administrative overhead.

upvoted 2 times

MatAlves 8 months, 3 weeks ago

"AWS Transit Gateway provides a hub and spoke design for connecting VPCs and on-premises networks as a fully managed service without requiring you to provision third-party virtual appliances. No VPN overlay is required, and AWS manages high availability and scalability."

<https://docs.aws.amazon.com/whitepapers/latest/building-scalable-secure-multi-vpc-network-infrastructure/transit-gateway.html#:~:text=AWS%20Transit%20Gateway%20provides%20a,manages%20high%20availability%20and%20scalability.>

upvoted 2 times

MatAlves 8 months, 3 weeks ago

Transit Gateway enables customers to connect thousands of VPCs. You can attach all your hybrid connectivity (VPN and Direct Connect connections) to a single gateway, consolidating and controlling your organization's entire AWS routing configuration in one place (refer to the following figure)

upvoted 2 times

ogerber 1 year, 3 months ago

Selected Answer: C

high number of accounts and VPC to connect to on prem _> exactly the transit gateway use case

upvoted 3 times

KZ06 1 year, 4 months ago

Hi,

Seems like after question 684, the discussion are quite less and seems recent comments. Are these new sets of questions updated?

Anyone having any idea around this?

upvoted 2 times

MatAlves 8 months, 3 weeks ago

Anytime I see "the number of VPCs will increase", I immediately look for "transit gateway" as the least administrative overhead.

upvoted 2 times

Cali182 1 year, 4 months ago

Selected Answer: C

vote for C

upvoted 3 times

EZforeverman 1 year, 4 months ago

I think its C. LEAST administrative overhead. D can work but AWS direct connection and VPC peering configure require too much administrative overhead

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #687

Topic 1

A company that uses AWS needs a solution to predict the resources needed for manufacturing processes each month. The solution must use historical values that are currently stored in an Amazon S3 bucket. The company has no machine learning (ML) experience and wants to use a managed service for the training and predictions.

Which combination of steps will meet these requirements? (Choose two.)

A. Deploy an Amazon SageMaker model. Create a SageMaker endpoint for inference. **Most Voted**

B. Use Amazon SageMaker to train a model by using the historical data in the S3 bucket. **Most Voted**

C. Configure an AWS Lambda function with a function URL that uses Amazon SageMaker endpoints to create predictions based on the inputs.

D. Configure an AWS Lambda function with a function URL that uses an Amazon Forecast predictor to create a prediction based on the inputs.

E. Train an Amazon Forecast predictor by using the historical data in the S3 bucket.

Correct Answer: AB

Community vote distribution

**Comments**

MatAlves **Highly Voted** 8 months, 3 weeks ago

Selected Answer: AB

"Amazon Forecast is no longer available to new customers. Existing customers of Amazon Forecast can continue to use the service as normal"

"After careful consideration, we have made the decision to close new customer access to Amazon Forecast, effective July 29, 2024."

This question will either be removed or reformulated to exclude Forecast as the service is no longer available to new customers.
upvoted 10 times

Salilgen 5 months, 1 week ago

I agree: this is probably an old questions, but I think answer was D,E

Otherwise you need three options: A,B and C (to create predictions)

upvoted 1 times

bujuman Highly Voted 1 year, 1 month ago

Selected Answer: DE

Because of these assertions

- The company has no machine learning (ML) experience
- The company wants to use a managed service

We could be tempted to go for SageMaker that is the core AWS managed service for ML purposes.

But, but, if we consider this valuable information:

- A company that uses AWS needs a solution to predict the resources needed for manufacturing processes.

With a bit of research, we will find out that AWS also holds time-series forecasting service based on machine learning (ML).

https://aws.amazon.com/forecast/?nc1=h_ls

So I understand options DE are the best answers even though this service is not mentioned anywhere in current SAA-C03 course version

upvoted 5 times

hilker1983 Most Recent 5 months, 1 week ago

Selected Answer: DE

D. Configure an AWS Lambda function with a function URL that uses an Amazon Forecast predictor to create a prediction based on the inputs.

Amazon Forecast is a managed service specifically designed for time series forecasting and can be used without extensive ML knowledge. This option allows for easy integration with Lambda for making predictions.

E. Train an Amazon Forecast predictor by using the historical data in the S3 bucket.

This option is essential for creating the forecasting model using the historical data stored in S3. Amazon Forecast is designed to handle this type of task and is user-friendly for those without ML experience.

upvoted 2 times

LeonSauveterre 5 months, 3 weeks ago

Selected Answer: DE

Amazon Forecast is no longer available to new customers as of July 29, 2024. However, for the AWS Certified Solutions Architect – Associate (SAA-C03) exam, Amazon Forecast may still appear as a valid service since the exam content often lags behind recent announcements.

Since the company has no ML expertise, A B C are all out, because they require prior SageMaker knowledge. That would make the answers: D & E.

But IRL (in real life, not a web service), you must consider SageMaker Canvas, because it "includes the key capabilities found in Forecast, including the ability to train an ensemble of forecasting models using both statistical and neural network algorithms. It creates the best model for your dataset by generating base models for each algorithm, evaluating their performance, and then combining the top-performing models into an ensemble".

upvoted 2 times

FlyingHawk 4 months, 2 weeks ago

We should reword the DE as below:

Train a time series forecasting model using Amazon SageMaker Canvas with the historical data in the S3 bucket.

SageMaker Canvas allows you to import data directly from Amazon S3, prepare the dataset, and train time series forecasting models through its intuitive interface.

Configure an AWS Lambda function with a function URL that uses the SageMaker Canvas model to create predictions based on the inputs.

After training the model, you can deploy it and use an AWS Lambda function to invoke the model endpoint for generating predictions programmatically.

upvoted 1 times

LeonSauveterre 5 months, 3 weeks ago

<https://aws.amazon.com/blogs/machine-learning/transition-your-amazon-forecast-usage-to-amazon-sagemaker-canvas/>

upvoted 2 times

XXXXXINN 8 months ago

so what about AB since the Forecast is no longer available to new customers?

so what about AB since the forecast is no longer available to new customers:

upvoted 3 times

JoeTromundo 8 months, 2 weeks ago

Selected Answer: AB

Amazon Forecast is no longer available to new customers.

<https://aws.amazon.com/blogs/machine-learning/transition-your-amazon-forecast-usage-to-amazon-sagemaker-canvas/>

upvoted 5 times

TwinSpark 1 year ago

Selected Answer: DE

Amazon forecast can be trained by using data from S3:

<https://docs.aws.amazon.com/forecast/latest/dg/getting-started.html>

upvoted 3 times

Hung23 1 year, 1 month ago

Selected Answer: BE

BE from CHATGPT

upvoted 1 times

lenotc 1 year, 2 months ago

Selected Answer: BE

SageMaker and Forecast can directly utilize data within an S3

B) E)

<https://aws.amazon.com/blogs/compute/build-workflows-for-amazon-forecast-with-aws-step-functions/>

<https://docs.aws.amazon.com/sagemaker/latest/dg/train-model.html>

upvoted 1 times

TheLaPlanta 1 year, 2 months ago

Selected Answer: AB

A + B dude

upvoted 2 times

Ravan 1 year, 3 months ago

Selected Answer: AB

Yes, exactly. Steps B and A together constitute a comprehensive solution:

- Step B involves using Amazon SageMaker to train a machine learning model using historical data stored in the S3 bucket.
- Step A involves deploying the trained model as a SageMaker endpoint, allowing for real-time inference on new data.

This combination leverages Amazon SageMaker's managed services for both training and inference, meeting the company's requirements efficiently.

upvoted 4 times

bodakrishna 1 year, 3 months ago

A & B:

B. Amazon SageMaker is a managed service that provides built-in algorithms and tools for training machine learning models. You can use SageMaker to train a model using historical data stored in an S3 bucket. This meets the requirement of utilizing a managed service for training the model without requiring machine learning experience.

A. Once the model is trained using SageMaker, you can deploy it by creating a SageMaker endpoint for inference. This endpoint allows you to make predictions based on new data, fulfilling the requirement of predicting resources needed for manufacturing processes each month.

upvoted 3 times

1Alpha1 1 year, 3 months ago

Selected Answer: DE

E: Amazon Forecast is a fully managed service that uses machine learning (ML) to generate highly accurate forecasts without requiring any prior ML experience. Forecast is applicable in a wide variety of use cases, including estimating product demand, energy demand, workforce planning, computing cloud infrastructure usage, traffic demand, supply chain optimization, and

financial planning.

D: Publish demand using AWS Lambda, AWS Step Functions, and Amazon CloudWatch Events rule to periodically (hourly) query the database and write the past X-months (count from the current timestamp) demand data into the source Amazon S3.

<https://aws.amazon.com/blogs/machine-learning/automating-your-amazon-forecast-workflow-with-lambda-step-functions-and-cloudwatch-events-rule/>

upvoted 4 times

Cali182 1 year, 4 months ago

Selected Answer: BD

B & D is the right choice

upvoted 2 times

anikolov 1 year, 4 months ago

Selected Answer: DE

My votes are for DE based on statement from AWS site:

"Alternatively, if you are looking for a fully managed service to deliver highly accurate forecasts, without writing code, we recommend checking out Amazon Forecast. Amazon Forecast is a time-series forecasting service based on machine learning (ML) and built for business metrics analysis."

<https://aws.amazon.com/blogs/machine-learning/deep-demand-forecasting-with-amazon-sagemaker/>

upvoted 3 times

jaswantn 1 year, 3 months ago

Why E?

upvoted 1 times

betttty 1 year, 4 months ago

Explanation:

Training the Model with SageMaker (Option B):

Use Amazon SageMaker to train a machine learning model based on historical data. SageMaker simplifies the process of training, deploying, and managing machine learning models.

Creating Predictions with Amazon Forecast (Option D):

Use Amazon Forecast to create a predictor based on historical data. Forecast is designed for time-series forecasting, making it suitable for predicting resources needed for manufacturing processes each month.

Combining SageMaker for training and Amazon Forecast for predictions provides a comprehensive solution, and AWS Lambda can be used to integrate these services into your workflow.

upvoted 5 times

JackyCCK 1 year, 2 months ago

combination of steps so it cannot be B,D.

B D is two different solution

upvoted 1 times

Andy_09 1 year, 4 months ago

BE looks correct

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #688

Topic 1

A company manages AWS accounts in AWS Organizations. AWS IAM Identity Center (AWS Single Sign-On) and AWS Control Tower are configured for the accounts. The company wants to manage multiple user permissions across all the accounts.

The permissions will be used by multiple IAM users and must be split between the developer and administrator teams. Each team requires different permissions. The company wants a solution that includes new users that are hired on both teams.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create individual users in IAM Identity Center for each account. Create separate developer and administrator groups in IAM Identity Center. Assign the users to the appropriate groups. Create a custom IAM policy for each group to set fine-grained permissions.
- B. Create individual users in IAM Identity Center for each account. Create separate developer and administrator groups in IAM Identity Center. Assign the users to the appropriate groups. Attach AWS managed IAM policies to each user as needed for fine-grained permissions.
- C. Create individual users in IAM Identity Center. Create new developer and administrator groups in IAM Identity Center. Create new permission sets that include the appropriate IAM policies for each group. Assign the new groups to the appropriate accounts. Assign the new permission sets to the new groups. When new users are hired, add them to the appropriate group. **Most Voted**
- D. Create individual users in IAM Identity Center. Create new permission sets that include the appropriate IAM policies for each user. Assign the users to the appropriate accounts. Grant additional IAM permissions to the users from within specific accounts. When new users are hired, add them to IAM Identity Center and assign them to the accounts.

Correct Answer: C

Community vote distribution

C (100%)

Comments

LeonSauveterre 5 months, 3 weeks ago

Selected Answer: C

I would:

1. Use AWS IAM Identity Center (SSO) to centrally manage user access across all accounts.
2. Create Permission Sets in Identity Center:
 - One for the Developer Team, granting access to resources with developer-specific permissions (AmazonEC2FullAccess or custom policies or something).
 - Another one for the Administrator Team, granting administrator-specific permissions (AdministratorAccess or something).
3. Sync users and groups from the corporate identity provider (or manage groups directly in IAM Identity Center) to automatically apply these permissions.
4. Assign these Permission Sets to groups of users in Identity Center rather than individuals, ensuring new users automatically inherit the appropriate permissions based on their group.

That's why you choose C, with the least operational overhead.

upvoted 1 times

xBUGx 8 months, 3 weeks ago

Selected Answer: C

C is least overhead

upvoted 1 times

1Alpha1 10 months ago

Selected Answer: C

Check out this one. https://www.youtube.com/watch?v=y_n9xN5mg1g

upvoted 1 times

Moon239 10 months ago

Selected Answer: C

<https://docs.aws.amazon.com/controltower/latest/userguide/sso.html>

upvoted 2 times

Cali182 10 months ago

Selected Answer: C

Correct is C

upvoted 2 times

Andy_09 10 months, 1 week ago

The correct answer should be C

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #689

Topic 1

A company wants to standardize its Amazon Elastic Block Store (Amazon EBS) volume encryption strategy. The company also wants to minimize the cost and configuration effort required to operate the volume encryption check.

Which solution will meet these requirements?

- A. Write API calls to describe the EBS volumes and to confirm the EBS volumes are encrypted. Use Amazon EventBridge to schedule an AWS Lambda function to run the API calls.
- B. Write API calls to describe the EBS volumes and to confirm the EBS volumes are encrypted. Run the API calls on an AWS Fargate task.
- C. Create an AWS Identity and Access Management (IAM) policy that requires the use of tags on EBS volumes. Use AWS Cost Explorer to display resources that are not properly tagged. Encrypt the untagged resources manually.
- D. Create an AWS Config rule for Amazon EBS to evaluate if a volume is encrypted and to flag the volume if it is not encrypted. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

MatAlves 8 months, 3 weeks ago

D: no-brainer.
upvoted 1 times

asdfcdsxdcf 1 year, 3 months ago

Selected Answer: D

D looks right
upvoted 3 times

bodakrishna 1 year, 3 months ago

AWS Config allows you to define rules to automatically check the configuration of AWS resources against desired configurations. By creating a custom AWS Config rule specifically for Amazon EBS volumes to evaluate if they are encrypted, you can ensure consistent encryption across all volumes. If a volume is found to be unencrypted, it can be flagged for further

action. This solution automates the process of encryption checking, minimizing manual effort and ensuring standardization across the environment. Additionally, AWS Config provides a cost-effective solution compared to continuously running scripts or tasks.

upvoted 3 times

mestule 1 year, 4 months ago

Selected Answer: D

AWS Config is a service that enables you to assess, audit, and evaluate the configurations of your AWS resources. It can check whether your resources comply with certain conditions (such as being encrypted), and it can flag or take action on resources that do not comply.

upvoted 4 times

betttty 1 year, 4 months ago

D :

you could use a managed rule to quickly start assessing whether your Amazon Elastic Block Store (Amazon EBS) volumes are encrypted or whether specific tags are applied to your resources.

https://docs.aws.amazon.com/config/latest/developerguide/evaluate-config_use-managed-rules.html

upvoted 4 times

Andy_09 1 year, 4 months ago

Correct answer is D

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #690

Topic 1

A company regularly uploads GB-sized files to Amazon S3. After the company uploads the files, the company uses a fleet of Amazon EC2 Spot Instances to transcode the file format. The company needs to scale throughput when the company uploads data from the on-premises data center to Amazon S3 and when the company downloads data from Amazon S3 to the EC2 instances.

Which solutions will meet these requirements? (Choose two.)

- A. Use the S3 bucket access point instead of accessing the S3 bucket directly.
- B. Upload the files into multiple S3 buckets.
- C. Use S3 multipart uploads. **Most Voted**
- D. Fetch multiple byte-ranges of an object in parallel. **Most Voted**
- E. Add a random prefix to each object when uploading the files.

Correct Answer: CD

Community vote distribution

CD (100%)

Comments

betttty **Highly Voted** 1 year, 4 months ago

CD
C: Increase the file upload throughput
D: increase the file download throughput
upvoted 10 times

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: CD

C: Upload: Multipart clear,
D: Download: You can fetch a byte-range from an object, transferring only the specified portion. You can use concurrent connections to Amazon S3 to fetch different byte ranges from within the same object. This helps you achieve higher aggregate throughput versus a single whole-object request.

A: S3 Access Points can be easily scaled, but are typically used to simplify data access for any AWS service or customer application that stores data in S3.

E: Prefixes: You can increase your read or write performance by using parallelization. For example, if you create 10 prefixes in an Amazon S3 bucket to parallelize reads, you could scale your read performance to 55,000 read requests per second.
But wording in this answer is strange...

upvoted 5 times

Salahdin Most Recent 6 months, 2 weeks ago

Selected Answer: CD

CD

C: Increase the file upload throughput

D: increase the file download throughpu

upvoted 2 times

dds69 1 year, 2 months ago

Selected Answer: CD

C&D are correct

upvoted 3 times

Bazzix 1 year, 2 months ago

Selected Answer: CD

Cd are correct

upvoted 3 times

bodakrishna 1 year, 3 months ago

C &D Correct

upvoted 3 times

Darshan07 1 year, 3 months ago

Selected Answer: CD

CD are the correct options

upvoted 3 times

Cali182 1 year, 4 months ago

Selected Answer: CD

CD is the correct for me

upvoted 2 times

Andy_09 1 year, 4 months ago

Correct answer is CD

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #691

Topic 1

A solutions architect is designing a shared storage solution for a web application that is deployed across multiple Availability Zones. The web application runs on Amazon EC2 instances that are in an Auto Scaling group. The company plans to make frequent changes to the content. The solution must have strong consistency in returning the new content as soon as the changes occur.

Which solutions meet these requirements? (Choose two.)

A. Use AWS Storage Gateway Volume Gateway Internet Small Computer Systems Interface (iSCSI) block storage that is mounted to the individual EC2 instances.

B. Create an Amazon Elastic File System (Amazon EFS) file system. Mount the EFS file system on the individual EC2 instances. **Most Voted**


C. Create a shared Amazon Elastic Block Store (Amazon EBS) volume. Mount the EBS volume on the individual EC2 instances.

D. Use AWS DataSync to perform continuous synchronization of data between EC2 hosts in the Auto Scaling group.

E. Create an Amazon S3 bucket to store the web content. Set the metadata for the Cache-Control header to no-cache. Use Amazon CloudFront to deliver the content. **Most Voted**

Correct Answer: BE

Community vote distribution



Comments

Andy_09 **Highly Voted** 1 year, 4 months ago

Correct answer BE
upvoted 8 times

Cali182 **Highly Voted** 1 year, 4 months ago

Selected Answer: BE

B & E seems to be the most logic
upvoted 5 times

LeanSolutions **Most Recent** 5 months, 2 weeks ago

Leonsauveterre most recent 5 months, 3 weeks ago

Selected Answer: BE

A - Volume Gateway is typically on-prem and optimized for hybrid setups, not for fully cloud-based solutions.
B - EFS provides a fully managed, shared file storage system that works across multiple AZs.
C - Even with Multi-Attach, sharing EBS volumes across AZs is pretty complex and error-prone.
D - DataSync is designed for periodic synchronization and not real-time, frequent updates.
E - The "no-cache" header ensures the content is always fetched fresh, and CloudFront delivers it efficiently with reduced latency. Although S3 is less efficient for frequent updates compared to a shared file system like EFS, but not a problem here.

upvoted 3 times

Penjerla 6 months ago

Selected Answer: BC

Option A is out as storage gateway is suited for hybrid cloud use cases where on-premises workloads need access to AWS cloud storage

upvoted 1 times

Salilgen 5 months, 1 week ago

C is wrong: you cannot share EBS in multiple AZ

upvoted 1 times

sandordini 1 year, 1 month ago

Regarding storage, I'd go for EFS, although it never mentions the requirement for file storage. Datasync can copy data between several storage types, including EFS, agents can be installed on EC2, but you cannot perform continuous synchronization of EC2 instances. Only storage. Cloudfront can publish both passive (s3) and active content (EC2+EFS) but wording doesn't tell a thing about such a share. And if it's a passive site why do we even have 2 storage types... I'd say, for me, the least bad solution seems to be B + E.

upvoted 2 times

Hung23 1 year, 2 months ago

Selected Answer: BE

I choose BE

upvoted 2 times

alawada 1 year, 2 months ago

BD looks most logical to me - continuous changes required an update via DataSync

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #692

Topic 1


A company is deploying an application in three AWS Regions using an Application Load Balancer. Amazon Route 53 will be used to distribute traffic between these Regions.

Which Route 53 configuration should a solutions architect use to provide the MOST high-performing experience?

- A. Create an A record with a latency policy. **Most Voted**
- B. Create an A record with a geolocation policy.
- C. Create a CNAME record with a failover policy.
- D. Create a CNAME record with a geoproximity policy.

Correct Answer: A

Community vote distribution



Comments

Kezuko **Highly Voted** 8 months, 3 weeks ago

Selected Answer: A

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html>
upvoted 7 times

h0ng97_spare_002 **Highly Voted** 8 months, 2 weeks ago

Selected Answer: A

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html>
upvoted 6 times

LeonSauveterre **Most Recent** 5 months, 3 weeks ago

Selected Answer: A

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html>

To provide the MOST high-performing experience, the Latency-based routing policy is the most appropriate. It ensures requests are routed to the AWS Region that provides the fastest response times based on real-world network conditions.

upvoted 1 times

sandordini 7 months, 2 weeks ago

Selected Answer: A

1. Given the chance >always use Alias over a Cname<
2. Latency-based routing is for user experience. (low latency)

Failover is for DR, Geolocation for local restrictions/rights/language/currency, and geo-proximity is a more complex, biased location-based routing, not part of the SA Associate exam.

upvoted 4 times

NSA_Poker 6 months ago

Alias?! "A record" is NOT an Alias record; it's an ADDRESS RECORD. There's a difference in between an address record (A record) & an Alias.

An Address record will map to 1 or more IP addresses.
An Alias record will map to another name like a CNAME does.

We eliminate C&D bc we need an IP address that will give us the best performance; we distribute traffic to a certain IP address based on policy.

geolocation policy is defined by where the request comes from.

latency policy is defined by how fast (performance) we can reply.

upvoted 3 times

mohammadthainat 8 months, 1 week ago

Selected Answer: A

Geoproximity Policy routing users to resources based on their geographic location, routing based on geographic location may not always be the absolute lowest latency.

latency-based routing prioritizes user experience.

upvoted 4 times

TruthWS 8 months, 2 weeks ago

A is true

upvoted 3 times

cedser8 9 months, 1 week ago

Selected Answer: D

The correct is D, the question says "using an Application Load Balancer" the ALB has a DNS name assigned not an IP. A type A record will only allow you to point to an IPv4. If I'm wrong, happy to be corrected.

upvoted 3 times

dkw2342 8 months, 2 weeks ago

Answer A is correct.

Route53 uses an internal record type called ALIAS, but from a DNS point of view it's still an A record.

Just try it yourself, create an ALB and create a DNS record in Route53. While you can technically use a CNAME (for subdomains, see below), the wizard will guide you to use an A ALIAS record, which also makes the most sense.

The problem with CNAME records is that it's not possible to create them at the root level of the domain. Let's say your domain is somedomain.com - you can't create a CNAME for the apex of the domain (mydomain.com), only for subdomains (subdomain.mydomain.com).

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/resource-record-sets-choosing-alias-non-alias.html>

upvoted 5 times

bodakrishna 9 months, 2 weeks ago

ChatGPT:

The most high-performing experience in this scenario would be achieved by using:

D. Create a CNAME record with a geoproximity policy.

Geoproximity routing allows you to route traffic based on the geographic location of your users and your resources. This would distribute traffic to the AWS Region that is closest to the user, optimizing performance by reducing latency. It's particularly

distribute traffic to the AWS region that is closest to the user, optimizing performance by reducing latency. It's particularly useful when deploying applications across multiple regions to ensure users are directed to the closest region, minimizing network latency and providing the best user experience.

upvoted 1 times

sandordini 7 months, 2 weeks ago

And, exactly, this is the reason why you should not rely on a LANGUAGE MODEL when you need a solution architect's advice .

upvoted 5 times

osmk 9 months, 3 weeks ago

A

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-latency.html>

upvoted 2 times

haci 9 months, 3 weeks ago

Selected Answer: A

Based on previous questions, I believe A is correct. Because; the closest geolocated server doesn't necessarily provide the best performance. Geolocated load balancing is mostly used for serving location-specific content.

upvoted 3 times

1Alpha1 10 months ago

Selected Answer: A

Q. What is Amazon Route 53's Latency Based Routing (LBR) feature?

LBR (Latency Based Routing) is a new feature for Amazon Route 53 that helps you improve your application's performance for a global audience. You can run applications in multiple AWS regions and Amazon Route 53, using dozens of edge locations worldwide, will route end users to the AWS region that provides the lowest latency.

<https://aws.amazon.com/route53/faqs/>

upvoted 3 times

Cali182 10 months ago

Selected Answer: B

Why would you use a CNAME record?? Most suitable seems to be option B

upvoted 1 times

osmk 9 months, 3 weeks ago

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-latency.html>

upvoted 2 times

Typewriter101 9 months, 1 week ago

Not really sure but ALBs do not have a static ip address they have domains assigned to them and also an Elastic ip can't be attached to an ALB. So mainly a cname would be preferred here.

upvoted 2 times

Typewriter101 9 months, 1 week ago

But generally speaking it's not a bad idea. But yes A record alias name can point to it. and i don't think it's B cause even if it's geolocation it may not always result in a high-performing exp.

upvoted 2 times

Andy_09 10 months ago

Sorry changing to B.

upvoted 1 times

Andy_09 10 months, 1 week ago

D looks correct.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #693

Topic 1

A company has a web application that includes an embedded NoSQL database. The application runs on Amazon EC2 instances behind an Application Load Balancer (ALB). The instances run in an Amazon EC2 Auto Scaling group in a single Availability Zone.

A recent increase in traffic requires the application to be highly available and for the database to be eventually consistent.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Replace the ALB with a Network Load Balancer. Maintain the embedded NoSQL database with its replication service on the EC2 instances.
- B. Replace the ALB with a Network Load Balancer. Migrate the embedded NoSQL database to Amazon DynamoDB by using AWS Database Migration Service (AWS DMS).
- C. Modify the Auto Scaling group to use EC2 instances across three Availability Zones. Maintain the embedded NoSQL database with its replication service on the EC2 instances.
- D. Modify the Auto Scaling group to use EC2 instances across three Availability Zones. Migrate the embedded NoSQL database to Amazon DynamoDB by using AWS Database Migration Service (AWS DMS). **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

LeonSauveterre 5 months, 3 weeks ago

Selected Answer: D

Application Load Balancer (ALB) is better suited for HTTP/HTTPS traffic. A & B are out.

Maintaining replication and failover for the embedded database across AZs is frustratingly complex and operationally intensive. Migration to DynamoDB requires initial effort but reduces long-term operational overhead. So D.

upvoted 2 times

freedafeng 10 months, 3 weeks ago

I honestly don't think you can use db Migration service to migrate an embedded db.

...memory, some things you can use as migration service to migrate an embedded database.
upvoted 1 times

HTHK 1 year, 1 month ago

DDDDDDDDDDDDDDDDDDDDDD

upvoted 1 times

TruthWS 1 year, 2 months ago

Option D: let focus on HA + Scaling

upvoted 1 times

Kezuko 1 year, 2 months ago

Selected Answer: D

ASG for application HA + DynamoDB Scale for HA

upvoted 4 times

rubiteb 1 year, 3 months ago

B as it's highly available and has less operational overhead than D.

upvoted 1 times

dkw2342 1 year, 2 months ago

ALB -> NLB makes no sense and solution lacks HA for the app layer.

upvoted 1 times

NayeraB 1 year, 3 months ago

But wouldn't migrating an embedded database to a new one introduce operational overhead now and in the future?

upvoted 1 times

MatAlves 8 months, 3 weeks ago

No, the very opposite:

"Amazon DynamoDB is a fully managed, serverless, key-value NoSQL database that runs high-performance applications at any scale."

upvoted 1 times

1Alpha1 1 year, 3 months ago

Selected Answer: D

DynamoDB + Modifying the Auto Scaling group

upvoted 2 times

Cali182 1 year, 4 months ago

Selected Answer: D

Dynamo DB presents more advantages, because it would need less administrative effort

upvoted 2 times

Andy_09 1 year, 4 months ago

The correct option should be D

upvoted 4 times