

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #694

Topic 1

A company is building a shopping application on AWS. The application offers a catalog that changes once each month and needs to scale with traffic volume. The company wants the lowest possible latency from the application. Data from each user's shopping cart needs to be highly available. User session data must be available even if the user is disconnected and reconnects.

What should a solutions architect do to ensure that the shopping cart data is preserved at all times?

- A. Configure an Application Load Balancer to enable the sticky sessions feature (session affinity) for access to the catalog in Amazon Aurora.
- B. Configure Amazon ElastiCache for Redis to cache catalog data from Amazon DynamoDB and shopping cart data from the user's session. **Most Voted**
- C. Configure Amazon OpenSearch Service to cache catalog data from Amazon DynamoDB and shopping cart data from the user's session.
- D. Configure an Amazon EC2 instance with Amazon Elastic Block Store (Amazon EBS) storage for the catalog and shopping cart. Configure automated snapshots.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Tanidanindo **Highly Voted** 1 year, 2 months ago

Selected Answer: B

• ElastiCache is a managed in-memory data store service that is well-suited for managing session data in a distributed architecture.

upvoted 5 times

LeonSauveterre **Most Recent** 5 months, 3 weeks ago

Selected Answer: B

User session data must be available even if the user is disconnected and reconnects. Choose ElastiCache for Redis. That remains true for so many similar questions.

upvoted 1 times

upvoted 1 times

RC6 12 months ago

Selected Answer: B

B looks correct

upvoted 2 times

asdfcdsxdc 1 year, 3 months ago

why not A?

upvoted 3 times

24b2e9e 11 months, 2 weeks ago

because data to store(cache) will be bigger in size

upvoted 2 times

knben 1 year, 3 months ago

Selected Answer: B

session data must be available even if the user is disconnected and reconnects -> ElastiCache for Redis

upvoted 2 times

1Alpha1 1 year, 3 months ago

Selected Answer: B

B: ELB <--> ASG <--> ElastiCache <--> DynamoDB

upvoted 3 times

Andy_09 1 year, 4 months ago

B looks correct

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #695

Topic 1

A company is building a microservices-based application that will be deployed on Amazon Elastic Kubernetes Service (Amazon EKS). The microservices will interact with each other. The company wants to ensure that the application is observable to identify performance issues in the future.

Which solution will meet these requirements?

- A. Configure the application to use Amazon ElastiCache to reduce the number of requests that are sent to the microservices.
- B. Configure Amazon CloudWatch Container Insights to collect metrics from the EKS clusters. Configure AWS X-Ray to trace the requests between the microservices. **Most Voted**
- C. Configure AWS CloudTrail to review the API calls. Build an Amazon QuickSight dashboard to observe the microservice interactions.
- D. Use AWS Trusted Advisor to understand the performance of the application.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Cali182 **Highly Voted** 10 months ago

Selected Answer: B

Option B

Amazon CloudWatch Container Insights: This service provides monitoring and troubleshooting capabilities for containerized applications. It collects and aggregates metrics, logs, and events from Amazon EKS clusters and containers. This helps in monitoring the performance and health of microservices.

upvoted 9 times

Andy_09 **Highly Voted** 10 months, 1 week ago

Correct answer is B

upvoted 6 times

LeonSauveterre **Most Recent** 5 months, 3 weeks ago

Selected Answer: B

A - ElastiCache can reduce latency but does not directly provide observability.
B - CloudWatch Container Insights provides detailed metrics and logs for containers. AWS X-Ray traces requests across microservices
C - CloudTrail is designed for security auditing and API call tracking, not detailed observability or performance monitoring.
D - Trusted Advisor focuses on high-level recommendations and does not provide real-time observability for application interactions or performance issues.

upvoted 1 times

Linuslin 6 months, 3 weeks ago

Selected Answer: B

B is the correct answer.

Use CloudWatch Container Insights to collect, aggregate, and summarize metrics and logs from your containerized applications and microservices.

https://docs.aws.amazon.com/zh_tw/AmazonCloudWatch/latest/monitoring/ContainerInsights.html

AWS X-Ray collects data about requests that your application serves, and it helps you view, filter, and gain insights into that data to identify issues and opportunities for optimization.

https://docs.aws.amazon.com/zh_tw/prescriptive-guidance/latest/logging-monitoring-for-application-owners/x-ray.html

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #696

Topic 1

A company needs to provide customers with secure access to its data. The company processes customer data and stores the results in an Amazon S3 bucket.


All the data is subject to strong regulations and security requirements. The data must be encrypted at rest. Each customer must be able to access only their data from their AWS account. Company employees must not be able to access the data.

Which solution will meet these requirements?

- A. Provision an AWS Certificate Manager (ACM) certificate for each customer. Encrypt the data client-side. In the private certificate policy, deny access to the certificate for all principals except an IAM role that the customer provides.
- B. Provision a separate AWS Key Management Service (AWS KMS) key for each customer. Encrypt the data server-side. In the S3 bucket policy, deny decryption of data for all principals except an IAM role that the customer provides.
- C. Provision a separate AWS Key Management Service (AWS KMS) key for each customer. Encrypt the data server-side. In each KMS key policy, deny decryption of data for all principals except an IAM role that the customer provides. **Most Voted**
- D. Provision an AWS Certificate Manager (ACM) certificate for each customer. Encrypt the data client-side. In the public certificate policy, deny access to the certificate for all principals except an IAM role that the customer provides.

Correct Answer: C

Community vote distribution



Comments

tch 3 months ago

Selected Answer: C

A key policy is a resource policy for an AWS KMS key. Key policies are the primary way to control access to KMS keys. Every KMS key must have exactly one key policy. The statements in the key policy determine who has permission to use the KMS key and how they can use it.

upvoted 1 times

tch 3 months ago

You can also use IAM policies and grants to control access to the KMS key, but every KMS key must have a key policy.

<https://docs.aws.amazon.com/kms/latest/developerguide/key-policies.html>

upvoted 1 times

LeonSauveterre 5 months, 3 weeks ago

Selected Answer: B

ACM certificates are for securing communication (e.g., HTTPS/SSL), not for encrypting data at rest in S3. A and D are out.

Between B and C, they both create a separate key for each customer, and both work. So which is more cumbersome, managing S3 policies or managing KMS key policies?

Of course the latter one. In B, S3 bucket policy centralizes access rules for all keys, while in C, policies are distributed across individual keys. That's why I decide to choose B.

upvoted 2 times

RC6 12 months ago

Selected Answer: C

C looks correct

upvoted 2 times

BBR01 1 year, 1 month ago

Selected Answer: C

Actually I think neither B or C is correctly worded. If talking about key policy, should be "Modify the key's policy to grant the IAM user permissions for the kms:GenerateDataKey and kms:Decrypt actions at minimum."

If talking about bucket policy, should be "Deny GetObjects of particular customer without condition kms key equals 1234abcd...."

upvoted 4 times

mohammadthainat 1 year, 2 months ago

Selected Answer: C

Encryption at rest --> KMS

Each customer must be able to access only their data --> KMS Key Policies

upvoted 4 times

LeonSauveterre 5 months, 3 weeks ago

Each customer must be able to access only their data --> Deny decryption of data for all principals except an IAM role that the customer provides. This doesn't necessarily have to be KMS key policies. B & C both work in this scenario.

upvoted 3 times

Neung983 1 year, 3 months ago

Selected Answer: B

B.

Here's why this option is the best fit:

Server-Side Encryption: Encrypting data server-side with KMS ensures encryption happens transparently within AWS, eliminating the need for complex client-side management and potential security risks associated with user-managed keys.

Customer-Specific Keys: Utilizing separate KMS keys for each customer provides granular access control and encryption isolation. Each customer can only decrypt their data using their specific KMS key.

S3 Bucket Policy: By denying decryption permissions for all principals except the dedicated customer IAM role in the S3 bucket policy, unauthorized access, even from company employees, is prevented. This aligns with the requirement of customer-specific data access.

upvoted 3 times

Cali182 1 year, 4 months ago

Selected Answer: C

Option C

From Chapt

Option A is incorrect because using ACM certificates is typically for establishing secure communication over HTTPS and doesn't directly relate to encrypting data at rest in S3.

Option B is incorrect because while it suggests using AWS KMS keys for encryption, it mentions using S3 bucket policies for access control, which would not be appropriate for controlling decryption permissions.

Option D is incorrect because it suggests using ACM certificates for client-side encryption, which is not typically used for

encrypting data at rest in S3, and the approach described would not effectively control access to the encrypted data.
upvoted 4 times

Andy_09 1 year, 4 months ago

Correct answer should be C
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #697

Topic 1


A solutions architect creates a VPC that includes two public subnets and two private subnets. A corporate security mandate requires the solutions architect to launch all Amazon EC2 instances in a private subnet. However, when the solutions architect launches an EC2 instance that runs a web server on ports 80 and 443 in a private subnet, no external internet traffic can connect to the server.

What should the solutions architect do to resolve this issue?

- A. Attach the EC2 instance to an Auto Scaling group in a private subnet. Ensure that the DNS record for the website resolves to the Auto Scaling group identifier.
- B. Provision an internet-facing Application Load Balancer (ALB) in a public subnet. Add the EC2 instance to the target group that is associated with the ALB. Ensure that the DNS record for the website resolves to the ALB. **Most Voted**
- C. Launch a NAT gateway in a private subnet. Update the route table for the private subnets to add a default route to the NAT gateway. Attach a public Elastic IP address to the NAT gateway.
- D. Ensure that the security group that is attached to the EC2 instance allows HTTP traffic on port 80 and HTTPS traffic on port 443. Ensure that the DNS record for the website resolves to the public IP address of the EC2 instance.

Correct Answer: B

Community vote distribution



Comments

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: B

Not A - Autoscaling Irrelevant

B - ALB, route tables for the public subnet with a route to Priv subnet

C - "NAT gateway" is "to allow [outbound] internet traffic", but this is about inbound traffic

D - Instances are in the private subnet, therefore it won't work from the public.

upvoted 6 times

anikolov **Highly Voted** 1 year, 4 months ago

Selected Answer: B

B: Provision an internet-facing Application Load Balancer (ALB) in a public subnet makes more sense

...between an internet-facing application and database (DB), and a private subnet makes more sense
upvoted 5 times

LeonSauveterre Most Recent 5 months, 3 weeks ago

Selected Answer: B

- A - Auto Scaling groups are for managing instance scaling.
- B - The internet-facing ALB is deployed in a public subnet and acts as a bridge between the internet and the private EC2 instance. Then, The ALB forwards HTTP/HTTPS traffic to the EC2 instance in the private subnet.
- C - A NAT gateway provides outbound internet access for private instances, but it does not solve inbound traffic issues.
- D - Private EC2 instances do not have public IP addresses at all.

upvoted 1 times

8621a7c 7 months, 2 weeks ago

Selected Answer: C

Is the question ask to solve the external connection?

upvoted 1 times

waldirisantos 1 year, 1 month ago

Why not "D"?

upvoted 1 times

boluwatito 1 year, 1 month ago

Selected Answer: D

nsure that the security group attached to the EC2 instance allows inbound traffic on ports 80 and 443 from the desired sources (e.g., any IP or specific IP ranges). This allows external internet traffic to reach the web server running on the EC2 instance

upvoted 1 times

GOTJ 4 months, 3 weeks ago

But the EC2s must be located in a private subnet, which is not direct internet accesible by definition

upvoted 1 times

TruthWS 1 year, 2 months ago

B - because ALB do it better NAT

upvoted 2 times

Cali182 1 year, 4 months ago

Selected Answer: C

Option C from Chatgt

upvoted 1 times

jaswantn 1 year, 3 months ago

NAT Gateway stays in public subnet, not in private subnet. So, C can't be.

upvoted 6 times

lenotc 1 year, 2 months ago

NAT Gateway outbound connections

upvoted 2 times

mestule 1 year, 4 months ago

Selected Answer: B

B makes most sense

upvoted 4 times

Andy_09 1 year, 4 months ago

Changing to option D

upvoted 1 times

Andy_09 1 year, 4 months ago

C should be the correct answer
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #698

Topic 1

A company is deploying a new application to Amazon Elastic Kubernetes Service (Amazon EKS) with an AWS Fargate cluster. The application needs a storage solution for data persistence. The solution must be highly available and fault tolerant. The solution also must be shared between multiple application containers.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create Amazon Elastic Block Store (Amazon EBS) volumes in the same Availability Zones where EKS worker nodes are placed. Register the volumes in a StorageClass object on an EKS cluster. Use EBS Multi-Attach to share the data between containers.
- B. Create an Amazon Elastic File System (Amazon EFS) file system. Register the file system in a StorageClass object on an EKS cluster. Use the same file system for all containers. **Most Voted**
- C. Create an Amazon Elastic Block Store (Amazon EBS) volume. Register the volume in a StorageClass object on an EKS cluster. Use the same volume for all containers.
- D. Create Amazon Elastic File System (Amazon EFS) file systems in the same Availability Zones where EKS worker nodes are placed. Register the file systems in a StorageClass object on an EKS cluster. Create an AWS Lambda function to synchronize the data between file systems.

Correct Answer: B

Community vote distribution

B (100%)

Comments

boluwatito **Highly Voted** 8 months ago

Selected Answer: B

Overall, Amazon EFS provides a highly available, fault-tolerant, and shared storage solution with minimal operational overhead, making it the ideal choice for persisting data in an Amazon EKS Fargate cluster.

upvoted 5 times

TruthWS **Most Recent** 8 months, 2 weeks ago

B bcs EBS only attach one EC2

upvoted 3 times

asdfcdsxdc 9 months, 1 week ago

Selected Answer: B

B looks correct
upvoted 3 times

Naveena_Devanga 9 months, 3 weeks ago

B, The solution also must be shared between multiple application containers so attaching to each container is not a practical solution.

upvoted 4 times

Marunio 9 months, 4 weeks ago

Selected Answer: B

B is correct answer because it is high available - EBS isn't HA for that so A isn't dealing with request.

upvoted 3 times

jaswantn 10 months ago

Option A... EBS with multi attach does not provide HA so option B is more appropriate.

upvoted 2 times

dkw2342 8 months, 2 weeks ago

It's just plain wrong. Not getting HA with EBS multi attach is really the least of your problems. Mounting a regular FS in read/write mode on more than one machine will cause data corruption. You'd need a clustered filesystem.

upvoted 2 times

Andy_09 10 months, 1 week ago

Correct answer is B

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #699

Topic 1

A company has an application that uses Docker containers in its local data center. The application runs on a container host that stores persistent data in a volume on the host. The container instances use the stored persistent data.

The company wants to move the application to a fully managed service because the company does not want to manage any servers or storage infrastructure.

Which solution will meet these requirements?

- A. Use Amazon Elastic Kubernetes Service (Amazon EKS) with self-managed nodes. Create an Amazon Elastic Block Store (Amazon EBS) volume attached to an Amazon EC2 instance. Use the EBS volume as a persistent volume mounted in the containers.
- B. Use Amazon Elastic Container Service (Amazon ECS) with an AWS Fargate launch type. Create an Amazon Elastic File System (Amazon EFS) volume. Add the EFS volume as a persistent storage volume mounted in the containers. **Most Voted**
- C. Use Amazon Elastic Container Service (Amazon ECS) with an AWS Fargate launch type. Create an Amazon S3 bucket. Map the S3 bucket as a persistent storage volume mounted in the containers.
- D. Use Amazon Elastic Container Service (Amazon ECS) with an Amazon EC2 launch type. Create an Amazon Elastic File System (Amazon EFS) volume. Add the EFS volume as a persistent storage volume mounted in the containers.

Correct Answer: B*Community vote distribution*

B (92%)

C (8%)

Comments**Marunio** **Highly Voted** 1 year, 3 months ago**Selected Answer: B**

Mounting S3 in Fargate is not supported commonly. You'd have to make it manually. EFS is very well supported with Fargate. <https://stackoverflow.com/questions/66391791/how-to-mount-s3-bucket-to-ecs-fargate-container>

<https://docs.aws.amazon.com/AmazonECS/latest/bestpracticesguide/storage.html>

upvoted 8 times

Andy_09 **Highly Voted** 1 year, 4 months ago

B looks correct
upvoted 6 times

waldirlsantos Most Recent 1 year, 1 month ago

Selected Answer: B

EFS is listed like a best practice for this cases

"ou can use Amazon ECS to run stateful containerized applications at scale by using AWS storage services, such as Amazon EFS, Amazon EBS, or FSx for Windows File Server, that provide data persistence to inherently ephemeral containers. The term data persistence means that the data itself outlasts the process that created it. "

upvoted 2 times

MattBJ 1 year, 2 months ago

Selected Answer: B

B is correct
upvoted 2 times

ogerber 1 year, 3 months ago

Selected Answer: C

The company does not want to manage any servers or storage infrastructure.
I would go with C

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Both B and C work with AWS launch type. So you have to decide between EFS vs S3.

How are you going to MOUNT s3 buckets in the containers?!

upvoted 2 times

MatAlves 8 months, 3 weeks ago

It's simply not good practice, unless you're confusing S3 with S3 storage gateway.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #700

Topic 1


A gaming company wants to launch a new internet-facing application in multiple AWS Regions. The application will use the TCP and UDP protocols for communication. The company needs to provide high availability and minimum latency for global users.

Which combination of actions should a solutions architect take to meet these requirements? (Choose two.)

- A. Create internal Network Load Balancers in front of the application in each Region. **Most Voted**
- B. Create external Application Load Balancers in front of the application in each Region.
- C. Create an AWS Global Accelerator accelerator to route traffic to the load balancers in each Region. **Most Voted**
- D. Configure Amazon Route 53 to use a geolocation routing policy to distribute the traffic.
- E. Configure Amazon CloudFront to handle the traffic and route requests to the application in each Region

Correct Answer: AC

Community vote distribution



Comments

Andy_09 **Highly Voted** 10 months, 1 week ago

Correct answer should be AC
upvoted 10 times

mestule 10 months ago

Agreed.

When you add an internal Load Balancer or an Amazon EC2 instance endpoint in AWS Global Accelerator, you enable internet traffic to flow directly to and from the endpoint in Virtual Private Clouds (VPCs) by targeting it in a private subnet. The VPC that contains the load balancer or EC2 instance must have an internet gateway attached to it, to indicate that the VPC accepts internet traffic. However, you don't need public IP addresses on the load balancer or EC2 instance. You also don't need an associated internet gateway route for the subnet.

upvoted 8 times

ogerber **Highly Voted** 9 months, 3 weeks ago

Selected Answer: AC

Gaming : TCP / UDP => always think NLB and global accelerator

Gaming + TCP / UDP => always think NLB and global accelerator

upvoted 5 times

yangbo Most Recent 3 months ago

Selected Answer: CD

A is wrong . AWS Global Accelerator cannot directly connect to Internal Network Load Balancers (NLBs). It can only connect to Internet-facing NLBs or Application Load Balancers (ALBs), and cannot connect to Internal NLBs because Internal NLBs are designed to serve traffic within the VPC and are not exposed to the internet.

upvoted 2 times

yangbo 3 months ago

client - dns - AWS Global Accelerator - NLB -EC2

upvoted 1 times

sandordini 7 months, 2 weeks ago

Selected Answer: AC

Gaming, TCP&UDP, HA, Low latency >> NLB + AWS Global Accelerator

upvoted 2 times

waldirisantos 7 months, 4 weeks ago

Selected Answer: AC

Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP. NLB + GA for UDP, TCP

upvoted 2 times

Kezuko 8 months, 3 weeks ago

Selected Answer: AC

UDP -> NLB and Global Accelerator

upvoted 4 times

1Alpha1 10 months ago

Selected Answer: AC

AC - the app is using TCP & UDP

upvoted 3 times

jaswantn 10 months ago

For global user where TCP and UDP protocols are used and HA with minimum latency is needed.... Global Accelerator with NLB is the solution combination .

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #701

Topic 1

A city has deployed a web application running on Amazon EC2 instances behind an Application Load Balancer (ALB). The application's users have reported sporadic performance, which appears to be related to DDoS attacks originating from random IP addresses. The city needs a solution that requires minimal configuration changes and provides an audit trail for the DDoS sources.

Which solution meets these requirements?

- A. Enable an AWS WAF web ACL on the ALB, and configure rules to block traffic from unknown sources.
- B. Subscribe to Amazon Inspector. Engage the AWS DDoS Response Team (DRT) to integrate mitigating controls into the service.
- C. Subscribe to AWS Shield Advanced. Engage the AWS DDoS Response Team (DRT) to integrate mitigating controls into the service. **Most Voted**
- D. Create an Amazon CloudFront distribution for the application, and set the ALB as the origin. Enable an AWS WAF web ACL on the distribution, and configure rules to block traffic from unknown sources

Correct Answer: C

Community vote distribution

C (100%)

Comments

Andy_09 **Highly Voted** 1 year, 4 months ago

C is the correct answer
upvoted 7 times

Scheldon **Most Recent** 11 months, 1 week ago

Selected Answer: C

AnswerC
upvoted 1 times

NSA_Poker 12 months ago

Selected Answer: C

(A & D) are incorrect.

AWS WAF Web ACL - contain WAF rules that define how to inspect web requests and what to do when a web request matches the inspection criteria.

We don't have the inspection criteria necessary to use WAF Web ACL effectively bc DDoS attacks are originating from random IP addresses.

The AWS DDoS Response Team can respond to the randomness.

(B) is incorrect.

Amazon Inspector - a service that analyzes your EC2 instances to identify potential security and configuration issues.

Inspector is not good at dealing with an actual DDOS attack like AWS Shield Advanced.

upvoted 1 times

sandordini 1 year, 1 month ago

Selected Answer: C

DDoS = AWS Shield

upvoted 2 times

Mikado211 1 year, 2 months ago

Selected Answer: C

C is the correct answer, AWS Shield Advanced.

upvoted 1 times

asdfcdsxdcf 1 year, 3 months ago

Selected Answer: C

C looks correct

upvoted 1 times

Naveena_Devanga 1 year, 3 months ago

C is the correct answer.

Amazon Inspector is an automated vulnerability management service whereas AWS Shield Advanced is a managed service that helps you protect your application against external threats, like DDoS attacks, volumetric bots, and vulnerability exploitation attempts. For higher levels of protection against attacks.

upvoted 2 times

Darshan07 1 year, 3 months ago

Selected Answer: C

C is the correct answer

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #702

Topic 1


A company copies 200 TB of data from a recent ocean survey onto AWS Snowball Edge Storage Optimized devices. The company has a high performance computing (HPC) cluster that is hosted on AWS to look for oil and gas deposits. A solutions architect must provide the cluster with consistent sub-millisecond latency and high-throughput access to the data on the Snowball Edge Storage Optimized devices. The company is sending the devices back to AWS.

Which solution will meet these requirements?

- A. Create an Amazon S3 bucket. Import the data into the S3 bucket. Configure an AWS Storage Gateway file gateway to use the S3 bucket. Access the file gateway from the HPC cluster instances.
- B. Create an Amazon S3 bucket. Import the data into the S3 bucket. Configure an Amazon FSx for Lustre file system, and integrate it with the S3 bucket. Access the FSx for Lustre file system from the HPC cluster instances. **Most Voted**
- C. Create an Amazon S3 bucket and an Amazon Elastic File System (Amazon EFS) file system. Import the data into the S3 bucket. Copy the data from the S3 bucket to the EFS file system. Access the EFS file system from the HPC cluster instances.
- D. Create an Amazon FSx for Lustre file system. Import the data directly into the FSx for Lustre file system. Access the FSx for Lustre file system from the HPC cluster instances.

Correct Answer: B

Community vote distribution



Comments

Linuslin **Highly Voted** 1 year ago

Selected Answer: B

No direct integration between Snowball and FSx for Lustre. It must be via S3.
Snowball Edge (Storage Optimized) --> S3 --integrate--> FSx for Lustre

<https://docs.aws.amazon.com/fsx/latest/LustreGuide/create-dra-linked-data-repo.html>

<https://aws.amazon.com/tw/blogs/aws/enhanced-amazon-s3-integration-for-amazon-fsx-for-lustre/>
upvoted 15 times

Cali182 **Highly Voted** 1 year, 4 months ago

Selected Answer: D

Option D

Option A, B, and C involve using Amazon S3 or Amazon EFS as an intermediary storage layer, which may introduce additional latency and overhead, not meeting the requirement of consistent sub-millisecond latency. Therefore, Option D is the most suitable solution for this scenario.

upvoted 12 times

domper20232023 1 year, 3 months ago

The format on the Snowball device would be s3 compatible only. The FSx for Lustre file system can be created and then linked to the S3 bucket. The Lustre file system can then be mounted on the HPC workloads that need sub-millisecond latency to store data. Option B would be the correct option, assuming only S3 support on snowball.

upvoted 9 times

a8a1e0e **Most Recent** 3 months, 4 weeks ago

Selected Answer: B

Integration from Snowall to Fsc only with S3

upvoted 1 times

FlyingHawk 5 months, 1 week ago

Selected Answer: B

it is not possible to import data directly from Snowball Edge Storage Optimized devices to Amazon FSx for Lustre, FSx for Lustre is designed to work seamlessly with data stored in Amazon S3, providing integration for high-performance access and lazy loading.

Snowball Edge devices are designed to securely transfer data to AWS services, particularly Amazon S3.

<https://aws.amazon.com/blogs/storage/automatically-import-objects-from-amazon-s3-into-amazon-fsx-for-lustre/>

upvoted 1 times

Abdullah2004 9 months, 3 weeks ago

Selected Answer: D

The correct answer is D for sure

upvoted 1 times

MatAlves 8 months, 3 weeks ago

You cannot import data from snowball in any other destination other than S3. So no, D is INCORRECT. Don't get tricked just because they mentioned HPC.

upvoted 2 times

n999 10 months, 2 weeks ago

Selected Answer: B

B for sure

upvoted 4 times

DZRomero 11 months, 3 weeks ago

Selected Answer: D

Import data to AWS services: When AWS receives the device, the data is automatically imported into the designated AWS service or Amazon S3 bucket based on your configuration. For example, if you need to access the data from an HPC cluster running on AWS, you would import the data into an Amazon FSx for Lustre file system or Amazon S3, and then access it from your HPC cluster instances.

No need S3 bucket

upvoted 1 times

trinh_le 1 year, 1 month ago

Selected Answer: B

You cannot access the FSx for Lustre file system from the HPC cluster instances and this is only possible via S3

upvoted 5 times

sandordini 1 year, 1 month ago

Selected Answer: D

HPC = Lustre
upvoted 3 times

sandordini 1 year, 1 month ago

Extension: HPC= Lustre, but Snowball = S3, therefore: B
Sync from Snowball to S3 -> Link/integrate with Lustre

Correct answer: C
upvoted 2 times

sandordini 1 year, 1 month ago

Which is of course not C but B... :D Sorry...
So correct answer: :D
upvoted 2 times

sukjubae 1 year, 1 month ago

B is right
upvoted 3 times

alawada 1 year, 2 months ago

Selected Answer: D is right answer because it mentions sub-millisecond latency and high-throughput access
upvoted 1 times

mgrimandi 1 year, 2 months ago

B
<https://medium.com/@abylead/amazon-fsx-for-migration-and-certification-f3cb7b4dd843>
upvoted 2 times

MattBJ 1 year, 2 months ago

Selected Answer: B

B is correct
upvoted 3 times

[Removed] 1 year, 3 months ago

Selected Answer: B

According to Copilot: Transferring data directly from AWS Snowball Edge to Amazon FSx for Lustre is not a standard process supported directly by AWS.
upvoted 3 times

iczcezar 1 year, 3 months ago

Selected Answer: D

Option D, creating an Amazon FSx for Lustre file system and importing the data directly into it, is indeed the most suitable solution for this scenario. By bypassing an intermediary storage layer and directly importing the data into FSx for Lustre, the solution ensures optimal performance with consistent sub-millisecond latency and high throughput, meeting the requirements of the HPC cluster. Thank you for pointing out the clarity.
upvoted 1 times

FZA24 1 year, 3 months ago

Selected Answer: B

It should be B.
No direct integration between Snowball and FSx for Lustre
upvoted 4 times

FZA24 1 year, 3 months ago

It must be via S3
upvoted 2 times

67a3f40 1 year, 3 months ago

07/03/19 1 year, 8 months ago

Cali182 you cannot directly copy from Snowball Edge to FSx for luster
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #703

Topic 1

A company has NFS servers in an on-premises data center that need to periodically back up small amounts of data to Amazon S3.

Which solution meets these requirements and is MOST cost-effective?

- A. Set up AWS Glue to copy the data from the on-premises servers to Amazon S3.
- B. Set up an AWS DataSync agent on the on-premises servers, and sync the data to Amazon S3. **Most Voted**
- C. Set up an SFTP sync using AWS Transfer for SFTP to sync data from on premises to Amazon S3.
- D. Set up an AWS Direct Connect connection between the on-premises data center and a VPC, and copy the data to Amazon S3.

Correct Answer: B

Community vote distribution

B (100%)

Comments

BillaRanga **Highly Voted** 1 year, 3 months ago

Selected Answer: B

A -> Used for ETL not copying

B -> Works

C -> Works, but overkill for the described scenario of periodic small backups, high cost

D -> Works but it may not be necessary for transferring small amounts of data periodically. High setup cost

upvoted 9 times

Andy_09 **Highly Voted** 1 year, 4 months ago

B is the correct option

upvoted 9 times

Scheldon **Most Recent** 11 months, 1 week ago

Selected Answer: B

AnswerB

Should be sufficient

upvoted 2 times

Darshan07 1 year, 3 months ago

Selected Answer: B

B is the correct option

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #704

Topic 1

An online video game company must maintain ultra-low latency for its game servers. The game servers run on Amazon EC2 instances. The company needs a solution that can handle millions of UDP internet traffic requests each second.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure an Application Load Balancer with the required protocol and ports for the internet traffic. Specify the EC2 instances as the targets.
- B. Configure a Gateway Load Balancer for the internet traffic. Specify the EC2 instances as the targets.
- C. Configure a Network Load Balancer with the required protocol and ports for the internet traffic. Specify the EC2 instances as the targets. **Most Voted**
- D. Launch an identical set of game servers on EC2 instances in separate AWS Regions. Route internet traffic to both sets of EC2 instances.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Andy_09** **Highly Voted** 1 year, 4 months ago

UDP needs NLB
upvoted 9 times

echonesis **Most Recent** 6 months, 3 weeks ago

UDP -> L4 Protocol -> NLB
upvoted 1 times

MatAlves 8 months, 3 weeks ago**Selected Answer: C**

UDP > NLB.
upvoted 2 times

zinabu 1 year, 2 months ago

Ans : C

OfCourse we can use both NLB and GLB balancers for UDP traffic but NLB is more cost effective than GLB that is why we choice C.

upvoted 1 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: C

TCP/UDP = NLB

upvoted 3 times

osmk 1 year, 3 months ago

C -><https://docs.aws.amazon.com/elasticloadbalancing/latest/network/introduction.html>

upvoted 4 times

Marunio 1 year, 3 months ago

Selected Answer: C

UDP -> NLB.

ALB is for HTTP/HTTPS.

Gateway Load Balancer is for 3rd party virtual appliances like Firewalls etc not the traffic distribution.

<https://aws.amazon.com/compare/the-difference-between-the-difference-between-application-network-and-gateway-load-balancing/#:~:text=An%20NLB%20operates%20on%20layer,level%20along%20with%20gateway%20functionality.>

upvoted 3 times

Gagg 1 year, 3 months ago

Selected Answer: C

UDP, should use network load balancer

upvoted 2 times

nj1999 1 year, 4 months ago

C, NLB

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #705

Topic 1

A company runs a three-tier application in a VPC. The database tier uses an Amazon RDS for MySQL DB instance.


The company plans to migrate the RDS for MySQL DB instance to an Amazon Aurora PostgreSQL DB cluster. The company needs a solution that replicates the data changes that happen during the migration to the new database.

Which combination of steps will meet these requirements? (Choose two.)

- A. Use AWS Database Migration Service (AWS DMS) Schema Conversion to transform the database objects. **Most Voted**
- B. Use AWS Database Migration Service (AWS DMS) Schema Conversion to create an Aurora PostgreSQL read replica on the RDS for MySQL DB instance.
- C. Configure an Aurora MySQL read replica for the RDS for MySQL DB instance.
- D. Define an AWS Database Migration Service (AWS DMS) task with change data capture (CDC) to migrate the data. **Most Voted**
- E. Promote the Aurora PostgreSQL read replica to a standalone Aurora PostgreSQL DB cluster when the replica lag is zero.

Correct Answer: AD

Community vote distribution



Comments

h0ng97_spare_002 **Highly Voted** 1 year, 2 months ago

Selected Answer: AD

A: Correct. because need convert from MySQL to PostgreSQL

B: Wrong. Schema Conversion does not create an Aurora read replica

C: Wrong. Company wants to migrate to Aurora PostgreSQL, not Aurora MySQL

D: Correct. CDC task helps to capture ongoing change from source data store

E: Wrong. Although using Aurora Read Replica is an option for DB migration within the same Region, this question is asking for "combination of steps", which this option does not have another compatible option to pair with

The correct answer is "AD"

I heretore, answer is "AD"

upvoted 8 times

mestule **Highly Voted** 1 year, 4 months ago

AD makes sense to me, but I am not sure if that's the best answer.

upvoted 6 times

Andy_09 1 year, 4 months ago

Agreed. AD makes more sense !!

upvoted 4 times

Scheldon **Most Recent** 11 months, 1 week ago

Selected Answer: AD

AnswerAD

upvoted 2 times

xBUGx 1 year, 2 months ago

Lag many never be zero, then it will ne er be promoted to primary

upvoted 2 times

haci 1 year, 3 months ago

Selected Answer: AD

It's quite similar with Q.235, based on that discussion A-D makes more sense.

upvoted 4 times

1e22522 10 months ago

of course, sin mas

upvoted 1 times

Andy_09 1 year, 4 months ago

Correct answer BE

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #706

Topic 1

A company hosts a database that runs on an Amazon RDS instance that is deployed to multiple Availability Zones. The company periodically runs a script against the database to report new entries that are added to the database. The script that runs against the database negatively affects the performance of a critical application. The company needs to improve application performance with minimal costs.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Add functionality to the script to identify the instance that has the fewest active connections. Configure the script to read from that instance to report the total new entries.
- B. Create a read replica of the database. Configure the script to query only the read replica to report the total new entries.
- C. Instruct the development team to manually export the new entries for the day in the database at the end of each day.
- D. Use Amazon ElastiCache to cache the common queries that the script runs against the database.

Most Voted

Correct Answer: B

Community vote distribution

B (100%)

Comments

Moon239 **Highly Voted** 1 year, 4 months ago

Selected Answer: B

Read replica
upvoted 5 times

Scheldon **Most Recent** 11 months, 1 week ago

Selected Answer: B

Answer B
upvoted 2 times

giovanna_mag 1 year, 3 months ago

Selected Answer: B

B, read replica
upvoted 3 times

mestule 1 year, 4 months ago

Selected Answer: B

B looks correct
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #707

Topic 1

A company is using an Application Load Balancer (ALB) to present its application to the internet. The company finds abnormal traffic access patterns across the application. A solutions architect needs to improve visibility into the infrastructure to help the company understand these abnormalities better.

What is the MOST operationally efficient solution that meets these requirements?

- A. Create a table in Amazon Athena for AWS CloudTrail logs. Create a query for the relevant information.
- B. Enable ALB access logging to Amazon S3. Create a table in Amazon Athena, and query the logs. **Most Voted**
- C. Enable ALB access logging to Amazon S3. Open each file in a text editor, and search each line for the relevant information.
- D. Use Amazon EMR on a dedicated Amazon EC2 instance to directly query the ALB to acquire traffic access log information.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Andy_09 **Highly Voted** 1 year, 4 months ago

B is the correct answer
upvoted 11 times

Marunio **Highly Voted** 1 year, 3 months ago

Selected Answer: B

A - Cloudtrail is for API Calls and changes on AWS account.
B - Going for athena in S3. - Correct
C - Manual work
D - Distractor
upvoted 9 times

Salilgen **Most Recent** 5 months, 1 week ago

Selected Answer: B

<https://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-access-logs.html>

upvoted 1 times

Scheldon 11 months, 1 week ago

Selected Answer: B

Answer B

upvoted 2 times

[Removed] 1 year, 1 month ago

Correct answer is B

upvoted 2 times

Naveena_Devanga 1 year, 3 months ago

B -

Amazon Athena is an interactive query service that makes it easy to analyze data directly in Amazon Simple Storage Service (Amazon S3) using standard SQL.

upvoted 5 times

c48b4e2 1 year, 3 months ago

Why there is a "Correct answer" (the green bordered one) at all while most of the time the community thinks (correctly) otherwise?

upvoted 2 times

betttty 1 year, 3 months ago

why not A?

upvoted 3 times

Kezuko 1 year, 2 months ago

Access logs is an optional feature of Elastic Load Balancing that is disabled by default

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #708

Topic 1

A company wants to use NAT gateways in its AWS environment. The company's Amazon EC2 instances in private subnets must be able to connect to the public internet through the NAT gateways.

Which solution will meet these requirements?

- A. Create public NAT gateways in the same private subnets as the EC2 instances.
- B. Create private NAT gateways in the same private subnets as the EC2 instances.
- C. Create public NAT gateways in public subnets in the same VPCs as the EC2 instances. **Most Voted**
- D. Create private NAT gateways in public subnets in the same VPCs as the EC2 instances.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**anikolov** **Highly Voted** 10 months ago**Selected Answer: C**

Should be C: Public NAT GW in Public Subnet to have access to internet. Private NAT GW is used for VPC or on-prem
upvoted 16 times

mestule **Highly Voted** 10 months ago**Selected Answer: C**

I think the correct is C, because D would require more than just private NAT gateway.

Private – Instances in private subnets can connect to other VPCs or your on-premises network through a private NAT gateway. You can route traffic from the NAT gateway through a transit gateway or a virtual private gateway. You cannot associate an elastic IP address with a private NAT gateway. You can attach an internet gateway to a VPC with a private NAT gateway, but if you route traffic from the private NAT gateway to the internet gateway, the internet gateway drops the traffic.

<https://docs.aws.amazon.com/vpc/latest/userguide/vpc-nat-gateway.html>

upvoted 9 times

Kezuko **Most Recent** 8 months, 3 weeks ago**Selected Answer: C**

Public NAT Gateway in public subnets for the internet access
<https://docs.aws.amazon.com/vpc/latest/userguide/vpc-nat-gateway.html>
upvoted 3 times

knben 9 months, 2 weeks ago

Selected Answer: C

Public NAT GW in Public Subnet to have access to internet
upvoted 2 times

Andy_09 10 months, 1 week ago

Looks correct
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #709

Topic 1

A company has an organization in AWS Organizations. The company runs Amazon EC2 instances across four AWS accounts in the root organizational unit (OU). There are three nonproduction accounts and one production account. The company wants to prohibit users from launching EC2 instances of a certain size in the nonproduction accounts. The company has created a service control policy (SCP) to deny access to launch instances that use the prohibited types.

Which solutions to deploy the SCP will meet these requirements? (Choose two.)

A. Attach the SCP to the root OU for the organization.

B. Attach the SCP to the three nonproduction Organizations member accounts. **Most Voted**


C. Attach the SCP to the Organizations management account.

D. Create an OU for the production account. Attach the SCP to the OU. Move the production member account into the new OU.

E. Create an OU for the required accounts. Attach the SCP to the OU. Move the nonproduction member accounts into the new OU. **Most Voted**

Correct Answer: BE

Community vote distribution



Comments

anikolov **Highly Voted** 1 year, 4 months ago

Selected Answer: BE

My vote is for BE
upvoted 13 times

mestule **Highly Voted** 1 year, 4 months ago

Selected Answer: BE

I think it's B (directly attach) and E (attach via OU).
upvoted 6 times

Dz1990 **Most Recent** 2 weeks, 3 days ago

Selected Answer: BE

Selected Answer: BE

□ Approach 1: Direct Account Attachment (Option B)
SCP → Directly attached to each nonproduction account
Pros: Simple, direct, immediate
Cons: Must manage policy on multiple accounts
□ Approach 2: Organizational Unit Structure (Option E)
SCP → Nonproduction OU → All nonproduction accounts inherit
Pros: Cleaner organization, easier to manage, scalable
Cons: Requires OU restructuring
upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: BE

B - Attach the SPC to the three accounts
E - Creates an OU > moves the member accounts to OU > attach the SCP to OU

"If you apply an authorization policy (for example, a service control policy (SCP)), to the root, it applies to all organizational units (OUs) and member accounts in the organization."

"A" would also affect the one production account, which we clearly don't want.

You can "attach an SCP to a root, OU, or account"
upvoted 2 times

MatAlves 8 months, 3 weeks ago

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps.html

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_getting-started_concepts.html
upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: BE

Only the non-prods need to be limited.
upvoted 4 times

67a3f49 1 year, 3 months ago

According to GPT-4 it's AE:
A. Attach the SCP to the root OU for the organization. This approach will apply the SCP to all accounts under the organization, including both nonproduction and production accounts. However, without additional context or actions, this does not meet the requirement to exclude the production account from the restrictions.

E. Create an OU for the required accounts. Attach the SCP to the OU. Move the nonproduction member accounts into the new OU. This is the correct approach as it directly addresses the requirement. By creating a separate OU for nonproduction accounts and attaching the SCP to this OU, you can specifically target the policy to only those accounts, effectively exempting the production account from the restrictions.
upvoted 1 times

1Alpha1 1 year, 3 months ago

Selected Answer: AC

AC - same answer

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_getting-started_concepts.html
upvoted 1 times

MatAlves 8 months, 3 weeks ago

The link you provided says:

"If you apply an authorization policy (for example, a service control policy (SCP)), to the root, it applies to all organizational units (OUs) and member accounts in the organization."

"A" would also affect the one production account, which we clearly don't want.

You can "attach an SCP to a root, OU, or account"

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps.html

upvoted 2 times

Cali182 1 year, 4 months ago

Selected Answer: AD

From Chat

A. Attach the SCP to the root OU for the organization: Attaching the SCP to the root OU ensures that it applies to all member accounts within the organization, including both nonproduction and production accounts.

D. Create an OU for the production account. Attach the SCP to the OU. Move the production member account into the new OU: By creating a separate OU for the production account and attaching the SCP to that OU, you can ensure that the SCP only affects the nonproduction accounts while allowing the production account to operate without restrictions.

upvoted 2 times

Andy_09 1 year, 4 months ago

CE should be the correct answer

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #710

Topic 1

A company's website hosted on Amazon EC2 instances processes classified data stored in Amazon S3. Due to security concerns, the company requires a private and secure connection between its EC2 resources and Amazon S3.

Which solution meets these requirements?

- A. Set up S3 bucket policies to allow access from a VPC endpoint. **Most Voted**
- B. Set up an IAM policy to grant read-write access to the S3 bucket.
- C. Set up a NAT gateway to access resources outside the private subnet.
- D. Set up an access key ID and a secret access key to access the S3 bucket.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Ashy1313 **Highly Voted** 10 months ago

Selected Answer: A

A VPC endpoint enables customers to privately connect to supported AWS services .
upvoted 7 times

sandordini **Highly Voted** 7 months, 2 weeks ago

Selected Answer: A

I think this question asks about the connection not about authorization, and for a secure S3 connection (e.g. without internet exposure, etc.) should be a VPC endpoint.
upvoted 7 times

Naveena_Devanga **Most Recent** 9 months, 3 weeks ago

D is the correct answer.
upvoted 1 times

Darshan07 9 months, 4 weeks ago

Selected Answer: A

A is the correct answer
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #711

Topic 1

An ecommerce company runs its application on AWS. The application uses an Amazon Aurora PostgreSQL cluster in Multi-AZ mode for the underlying database. During a recent promotional campaign, the application experienced heavy read load and write load. Users experienced timeout issues when they attempted to access the application.

A solutions architect needs to make the application architecture more scalable and highly available.

Which solution will meet these requirements with the LEAST downtime?

- A. Create an Amazon EventBridge rule that has the Aurora cluster as a source. Create an AWS Lambda function to log the state change events of the Aurora cluster. Add the Lambda function as a target for the EventBridge rule. Add additional reader nodes to fail over to.
- B. Modify the Aurora cluster and activate the zero-downtime restart (ZDR) feature. Use Database Activity Streams on the cluster to track the cluster status.
- C. Add additional reader instances to the Aurora cluster. Create an Amazon RDS Proxy target group for the Aurora cluster.
Most Voted
- D. Create an Amazon ElastiCache for Redis cache. Replicate data from the Aurora cluster to Redis by using AWS Database Migration Service (AWS DMS) with a write-around approach.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Marunio **Highly Voted** 1 year, 3 months ago

Selected Answer: C

Only C is real viable option - Adding Reader replica for handling Read load and RDS Proxy for connections.
upvoted 7 times

Andy_09 **Highly Voted** 1 year, 4 months ago

I would go for option C
upvoted 5 times

sk1974 Most Recent 3 months, 1 week ago

Selected Answer: C

C is the most applicable among the 4 available options.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: C

C - Explanation below:

upvoted 2 times

MatAlves 8 months, 3 weeks ago

A - "Lambda function to log state changes" - doesn't help with read/write load.

B - ZDR applies to restarts that Aurora performs automatically to resolve error conditions: doesn't help with read/write load.

D - Write-around approach: data is always written to the database and the data that is read goes to the cache. Doesn't help with read/write load.

C - CORRECT. Even though it doesn't address "write operations", Aurora Replicas to offload read workloads from the primary DB instance. Amazon RDS Proxy allows applications to pool and share connections established with the database, improving database efficiency and application scalability.

upvoted 4 times

Scheldon 11 months, 2 weeks ago

Selected Answer: C

AnswerC.

Proxy should help with the problem

upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: C

RDX proxy to handle timeout issue

upvoted 2 times

xBUGx 1 year, 2 months ago

Selected Answer: C

I go with C bc there is no better option

upvoted 4 times

jaswantn 1 year, 3 months ago

RDX proxy to handle timeout issue. option C

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #712

Topic 1

A company is designing a web application on AWS. The application will use a VPN connection between the company's existing data centers and the company's VPCs.

The company uses Amazon Route 53 as its DNS service. The application must use private DNS records to communicate with the on-premises services from a VPC.

Which solution will meet these requirements in the MOST secure manner?

A. Create a Route 53 Resolver outbound endpoint. Create a resolver rule. Associate the resolver rule with the VPC.

Most Voted

B. Create a Route 53 Resolver inbound endpoint. Create a resolver rule. Associate the resolver rule with the VPC.

C. Create a Route 53 private hosted zone. Associate the private hosted zone with the VPC.

D. Create a Route 53 public hosted zone. Create a record for each service to allow service communication

Correct Answer: A*Community vote distribution*

A (88%)

C (12%)

Comments**haci** **Highly Voted** 1 year, 3 months ago**Selected Answer: A**

If you have workloads that leverage both VPCs and on-premises resources, you also need to resolve DNS records hosted on-premises. Similarly, these on-premises resources may need to resolve names hosted on AWS. Through Resolver endpoints and conditional forwarding rules, you can resolve DNS queries between your on-premises resources and VPCs to create a hybrid cloud setup over VPN or Direct Connect (DX). Specifically:

Inbound Resolver endpoints allow DNS queries to your VPC from your on-premises network or another VPC.

Outbound Resolver endpoints allow DNS queries from your VPC to your on-premises network or another VPC.

Reference: <https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/resolver.html>

upvoted 17 times

MatAlves 8 months, 3 weeks ago

Right to the point!

upvoted 2 times

JCVDB23 Highly Voted 1 year, 2 months ago

Selected Answer: A

Amazon Route 53 Resolver provides DNS resolution for VPCs and on-premises networks over a Direct Connect or VPN connection. An outbound resolver endpoint forwards DNS queries from your VPC to your on-premises DNS service. A resolver rule specifies the domain names for the DNS queries that you want to forward (such as example.com), and the IP addresses of the DNS resolvers in your on-premises network.

Option C is not suitable because private hosted zones are used to route traffic within a VPC

<https://aws.amazon.com/blogs/architecture/using-route-53-private-hosted-zones-for-cross-account-multi-region-architectures/>

upvoted 6 times

Joshua721 Most Recent 4 months, 3 weeks ago

Selected Answer: A

Inbound Resolver endpoints allow DNS queries to your VPC from your on-premises network or another VPC.

Outbound Resolver endpoints allow DNS queries from your VPC to your on-premises network or another VPC.

Resolver rules enable you to create one forwarding rule for each domain name and specify the name of the domain for which you want to forward DNS queries from your VPC to an on-premises DNS resolver and from your on-premises to your VPC. Rules are applied directly to your VPC and can be shared across multiple accounts.

upvoted 2 times

Duddle88 5 months, 2 weeks ago

Selected Answer: C

"The company uses Amazon Route 53 as its DNS service" THERE IS NO DNS ON-PREMISES

Private DNS Resolution: It allows you to create DNS records that are only accessible within your VPC, ensuring secure and private communication between your AWS resources and on-premises network.

Forwarding DNS Queries: This is useful if you need to forward DNS queries from your VPC to your on-premises DNS servers. It's typically used when you have existing DNS infrastructure on-premises that you want to continue using.

upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: A

AWS <-> On-premises = Route 53 Resolver

- Outbound Resolver = From your VPC (AWS) to On-premises or another VPC

- Inbound Resolver = From on-premises network or another VPC TO your VPC.

upvoted 3 times

Jacky_S 11 months, 2 weeks ago

Selected Answer: C

The reason why i vote on C, because the question mentioned that "The company uses Amazon Route53 as it's DNS service" and did not mention that is using multiple accounts, so it should be the most secure way to just add the record in it's private host zone of it's own account due to dns poisoning concern.

Of cause, i totally agree on A if the dns zone owner is in on-premises dns server which reduce the operation efforts.

upvoted 3 times

cjace 11 months, 4 weeks ago

C. Create a Route 53 private hosted zone. Associate the private hosted zone with the VPC.

This setup allows the application within the VPC to resolve DNS queries using private DNS records, ensuring that the communication remains within the AWS network and is not exposed to the public internet. Associating the private hosted zone with the VPC ensures that only the resources within the VPC can resolve the DNS queries, maintaining a secure environment for application and on-premises service communication.

The outbound resolver endpoint and rule would be more relevant if the requirement was for resources within the VPC to resolve DNS queries for domain names that are located in the on-premises network. In that case, the outbound resolver would forward queries from the VPC to the on-premises DNS server for resolution. However, for private DNS communication from the VPC to

queries from the VPC to the on-premises DNS server for resolution. However, for private DNS communication from the VPC to on-premises services, the private hosted zone is the most secure method.

upvoted 3 times

alawada 1 year, 2 months ago

Selected Answer: A

Amazon Route 53 Resolver provides DNS resolution for VPCs and on-premises networks

upvoted 2 times

anikolov 1 year, 4 months ago

Selected Answer: A

Should be A "Create a Route 53 Resolver outbound endpoint."

upvoted 5 times

Andy_09 1 year, 4 months ago

Looks correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #713

Topic 1

A company is running a photo hosting service in the us-east-1 Region. The service enables users across multiple countries to upload and view photos. Some photos are heavily viewed for months, and others are viewed for less than a week. The application allows uploads of up to 20 MB for each photo. The service uses the photo metadata to determine which photos to display to each user.

Which solution provides the appropriate user access MOST cost-effectively?

- A. Store the photos in Amazon DynamoDB. Turn on DynamoDB Accelerator (DAX) to cache frequently viewed items.
- B. Store the photos in the Amazon S3 Intelligent-Tiering storage class. Store the photo metadata and its S3 location in DynamoDB. **Most Voted**
- C. Store the photos in the Amazon S3 Standard storage class. Set up an S3 Lifecycle policy to move photos older than 30 days to the S3 Standard-Infrequent Access (S3 Standard-IA) storage class. Use the object tags to keep track of metadata.
- D. Store the photos in the Amazon S3 Glacier storage class. Set up an S3 Lifecycle policy to move photos older than 30 days to the S3 Glacier Deep Archive storage class. Store the photo metadata and its S3 location in Amazon OpenSearch Service.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Typewriter101 **Highly Voted** 1 year, 3 months ago

Selected Answer: B

The Intelligent-Tiering storage class automatically moves objects between two access tiers (frequent access and infrequent access) based on their access patterns, which aligns well with the varying view frequencies of the photos. Storing metadata in DynamoDB allows for efficient querying and retrieval of photo metadata.

upvoted 10 times

Andy_09 **Highly Voted** 1 year, 4 months ago

B is the correct option

upvoted 10 times

Q. H. Khan **Most Recent** 11 months ago

Scneider most recent 11 months, 2 weeks ago

Selected Answer: B

AnswerB

Seem to be the best solution from provided
upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: B

Store the photos in the Amazon S3 Intelligent-Tiering = Unpredictable scenario
upvoted 3 times

Indrasis 1 year, 3 months ago

Correct option: B
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #714

Topic 1

A company runs a highly available web application on Amazon EC2 instances behind an Application Load Balancer. The company uses Amazon CloudWatch metrics.


As the traffic to the web application increases, some EC2 instances become overloaded with many outstanding requests. The CloudWatch metrics show that the number of requests processed and the time to receive the responses from some EC2 instances are both higher compared to other EC2 instances. The company does not want new requests to be forwarded to the EC2 instances that are already overloaded.

Which solution will meet these requirements?

- A. Use the round robin routing algorithm based on the RequestCountPerTarget and ActiveConnectionCount CloudWatch metrics.
- B. Use the least outstanding requests algorithm based on the RequestCountPerTarget and ActiveConnectionCount CloudWatch metrics. **Most Voted**
- C. Use the round robin routing algorithm based on the RequestCount and TargetResponseTime CloudWatch metrics.
- D. Use the least outstanding requests algorithm based on the RequestCount and TargetResponseTime CloudWatch metrics.

Correct Answer: B

Community vote distribution



Comments

h0ng97_spare_002 **Highly Voted** 1 year, 2 months ago

Selected Answer: B

Option B is correct because can use "RequestCountPerTarget" to identify the amount of requests for each EC2 instance. Then use "least outstanding requests algorithm" to route to targets with the lowest number of in progress requests.

Option D is wrong because "RequestCount" cannot identify the amount of requests for each EC2 instance. "RequestCount" is for the whole ALB.

upvoted 11 times

Tatai2015 1 year ago

<https://aws.amazon.com/about-aws/whats-new/2019/11/application-load-balancer-now-supports-least-outstanding-requests-algorithm-for-load-balancing-requests/>

upvoted 3 times

Andy_09 Highly Voted 1 year, 4 months ago

Option B would be the correct choice

upvoted 8 times

chest_jd Most Recent 1 month, 3 weeks ago

Selected Answer: D

D is correct.

<https://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-cloudwatch-metrics.html>

ActiveConnectionCount - The total number of concurrent TCP connections active from clients to the load balancer

RequestCountPerTarget - The average request count per target, in a target group

Both metrics do not show stage of every target.

RequestCount - The number of requests processed over IPv4 and IPv6. This metric is only incremented for requests where the load balancer node was able to choose a target.

We can say how many requests go to targets.

TargetResponseTime - The time elapsed, in seconds, after the request leaves the load balancer until the target starts to send the response headers.

We have metric to get statistics for every target/instance.

In fact Item D is correct

upvoted 1 times

sk1974 3 months, 1 week ago

Selected Answer: B

Actually, I now vote for B and not D

upvoted 1 times

sk1974 3 months, 1 week ago

Selected Answer: D

I will go with D since the question asks for the target response.

upvoted 1 times

FlyingHawk 5 months, 1 week ago

Selected Answer: B

D will be correct if change RequestCount to RequestCountPerTarget. RequestCount applies to the ALB rather than individual targets, making it unsuitable for identifying overloaded instances.

ActiveConnectionCount:

This metric indicates the current number of active TCP connections. It complements RequestCountPerTarget by helping identify instances that might be overwhelmed with connections, even if their request count seems moderate.

upvoted 1 times

Scheldon 11 months, 2 weeks ago

Selected Answer: B

AnswerB

upvoted 2 times

TruthWS 1 year, 2 months ago

Option B

upvoted 3 times

dkw2342 1 year, 2 months ago

IMO the correct answer is option D:

This is from an earlier version of the AWS documentation on ALB target groups - for some reason they removed this information in the current revision:

"Consider using least outstanding requests when the requests for your application vary in complexity or your targets vary in processing capability. Round robin is a good choice when the requests and targets are similar, or if you need to distribute requests equally among targets. You can compare the effect of round robin versus least outstanding requests using the following CloudWatch metrics: RequestCount, TargetConnectionErrorCount, and TargetResponseTime."

<https://web.archive.org/web/20200426172626/https://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-target-groups.html#modify-routing-algorithm>

upvoted 2 times

Salilgen 5 months, 1 week ago

I agree.

"The CloudWatch metrics show that the number of requests processed and the time to receive the responses from some EC2 instances are both higher compared to other EC2 instances". Then we must monitor time and number of requests metrics

upvoted 1 times

h0ng97_spare_002 1 year, 2 months ago

The link is just saying that you can view "RequestCount, TargetConnectionErrorCount, and TargetResponseTime" to understand the difference in effect between round robin vs least outstanding requests. It is not the direct answer to this question.

upvoted 3 times

xBUGx 1 year, 2 months ago

Selected Answer: D

I think TargetResponseTime is the best indicator for telling is a server is overloaded or not

upvoted 1 times

alawada 1 year, 2 months ago

Selected Answer: B

distribute the number of requests among instances

upvoted 2 times

Kezuko 1 year, 2 months ago

Selected Answer: B

<https://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-cloudwatch-metrics.html>

To understand the types

upvoted 4 times

haci 1 year, 3 months ago

Selected Answer: B

The question is not asking for better performance in response time. It is just simply asking to distribute the number of requests among instances. So B seems more logical.

upvoted 3 times

osmk 1 year, 3 months ago

Selected Answer: D

The least outstanding requests routing algorithm routes requests to the targets with the lowest number of in progress requests > <https://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-target-groups.html>

upvoted 3 times

osmk 1 year, 3 months ago

D>>> The least outstanding requests routing algorithm routes requests to the targets with the lowest number of in progress requests > <https://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-target-groups.html>

upvoted 1 times

Moon239 1 year, 4 months ago

Why not D?

upvoted 2 times

jaswantn 1 year, 3 months ago

With Least outstanding requests algorithm, new request will send it to the "target" with least number of outstanding requests. Targets processing long-standing requests or having lower processing capabilities are not burdened with more requests. That's why option B is correct and not option D.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #715

Topic 1

A company uses Amazon EC2, AWS Fargate, and AWS Lambda to run multiple workloads in the company's AWS account. The company wants to fully make use of its Compute Savings Plans. The company wants to receive notification when coverage of the Compute Savings Plans drops.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Create a daily budget for the Savings Plans by using AWS Budgets. Configure the budget with a coverage threshold to send notifications to the appropriate email message recipients. **Most Voted**
- B. Create a Lambda function that runs a coverage report against the Savings Plans. Use Amazon Simple Email Service (Amazon SES) to email the report to the appropriate email message recipients.
- C. Create an AWS Budgets report for the Savings Plans budget. Set the frequency to daily.
- D. Create a Savings Plans alert subscription. Enable all notification options. Enter an email address to receive notifications.

Correct Answer: A*Community vote distribution***Comments****anikolov** **Highly Voted** 1 year, 4 months ago**Selected Answer: A**

My vote is for A : <https://docs.aws.amazon.com/savingsplans/latest/userguide/sp-usingBudgets.html>
upvoted 10 times

lenote **Highly Voted** 1 year, 2 months ago**Selected Answer: D**

D:
<https://aws.amazon.com/about-aws/whats-new/2020/11/savings-plans-alerts-now-available-in-aws-cost-management/>
upvoted 6 times

sk1974 **Most Recent** 3 months, 1 week ago**Selected Answer: D**

Question asks when 'The company wants to receive notification when coverage of the Compute Savings Plans drops'. So, I will

Question asks when the company wants to receive notification when coverage of the Compute Savings Plans drops. . . 30 . I will say answer id D

upvoted 1 times

zdi561 3 months, 4 weeks ago

Selected Answer: A

D is alert for expiration. A is for coverage drop.

upvoted 2 times

Mistwalker 4 months, 2 weeks ago

Selected Answer: D

It sure seems to be this:

<https://aws.amazon.com/about-aws/whats-new/2020/11/savings-plans-alerts-now-available-in-aws-cost-management/>

upvoted 1 times

rosanna 5 months, 3 weeks ago

Selected Answer: D

AWS budget is mostly used to avoid exceeding the threshold, not under utilization alert.

upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: A

A - describes exactly what is said in this link:

"You can use AWS Budgets to enable simple-to-complex cost and usage tracking. Some examples include:

(...)

Setting a daily utilization or coverage budget to track your RI or Savings Plans. You can choose to be notified through email and Amazon SNS topics when your utilization drops below 80 percent for a given day."

<https://docs.aws.amazon.com/cost-management/latest/userguide/budgets-managing-costs.html>

upvoted 4 times

MatAlves 8 months, 3 weeks ago

D - clearly not what the question is asking:

"You can track your Savings Plans expirations and upcoming queued Savings Plans in Cost Explorer. You can use Savings Plans alerts to receive advance email alerts 1, 7, 30, or 60 days before your Savings Plan expiration date, or in when a commitment is queued for purchase. These notifications also alert you on the expiration date, and can be sent to up to 10 email recipients."

"Savings Plans Expiration Alerts notify you as your existing Savings Plans approach expiration. These alerts can help you easily renew your Savings Plans ."

upvoted 3 times

MatAlves 8 months, 3 weeks ago

<https://docs.aws.amazon.com/savingsplans/latest/userguide/sp-alert.html>

upvoted 2 times

Scheldon 11 months, 2 weeks ago

Selected Answer: A

AnswerA

We can set Savings Plan in AWS Budgets which will notify us if utilization, coverage and costs will be not in set range.

<https://docs.aws.amazon.com/savingsplans/latest/userguide/sp-usingBudgets.html>

upvoted 3 times

Kezuko 1 year, 2 months ago

Selected Answer: D

<https://aws.amazon.com/blogs/aws-cloud-financial-management/launch-savings-plans-expiration-and-queued-alerts-now-available-in-aws-cost-management/>

upvoted 4 times

YGHUIWRHF1234 1 year, 2 months ago

Selected Answer: A

Correct answer is A

upvoted 2 times

xBUGx 1 year, 2 months ago

Selected Answer: A

A is precisely targeted

upvoted 3 times

ManishGup 1 year, 3 months ago

Ny vote going to D.

<https://aws.amazon.com/blogs/aws-cloud-financial-management/launch-savings-plans-expiration-and-queued-alerts-now-available-in-aws-cost-management/>

upvoted 1 times

Indrasis 1 year, 3 months ago

Selected Answer: A

A is correct

upvoted 2 times

jaswantn 1 year, 3 months ago

Option D...In the Savings Plans Overview page indicate how many days in advance you would like to receive Savings Plans Alerts for Plan's expiration and upcoming queued purchase notifications.

upvoted 3 times

Andy_09 1 year, 4 months ago

Option D

upvoted 3 times

hajra313 1 year, 3 months ago

alert subscription will notify u before ending saving plan

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #716

Topic 1

A company runs a real-time data ingestion solution on AWS. The solution consists of the most recent version of Amazon Managed Streaming for Apache Kafka (Amazon MSK). The solution is deployed in a VPC in private subnets across three Availability Zones.


A solutions architect needs to redesign the data ingestion solution to be publicly available over the internet. The data in transit must also be encrypted.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Configure public subnets in the existing VPC. Deploy an MSK cluster in the public subnets. Update the MSK cluster security settings to enable mutual TLS authentication. **Most Voted**
- B. Create a new VPC that has public subnets. Deploy an MSK cluster in the public subnets. Update the MSK cluster security settings to enable mutual TLS authentication.
- C. Deploy an Application Load Balancer (ALB) that uses private subnets. Configure an ALB security group inbound rule to allow inbound traffic from the VPC CIDR block for HTTPS protocol.
- D. Deploy a Network Load Balancer (NLB) that uses private subnets. Configure an NLB listener for HTTPS communication over the internet.

Correct Answer: A

Community vote distribution



Comments

haci **Highly Voted** 1 year, 3 months ago

Selected Answer: A

Since we are talking about real-time data (UDP packets) ALB is not a viable solution. You don't need to listen HTTPS, so D is eliminated. If you create a new VPC, you must create link between the old one and this is not mentioned in B. So It is A for me.
upvoted 11 times

Dantecito **Most Recent** 3 months, 3 weeks ago

Selected Answer: D

D. NLB forward traffic from the internet to private subnet where the cluster is. MSK use layer 4.
A. B. use public subnets exposing the cluster
C. Similar to D but MSK requires a Layer 4.

upvoted 1 times

FlyingHawk 4 months, 1 week ago

Selected Answer: A

For MSK to get public access, you must associate it with a public subnet, so rule out C and D, B is to create a new VPC, which is more complicated and more operations, so A is correct.

upvoted 1 times

FlyingHawk 4 months, 1 week ago

To turn on public access to an MSK Provisioned cluster, first ensure that the cluster meets all of the following conditions:

The subnets that are associated with the cluster must be public. Each public subnet has a public IPv4 address associated with it and public IPv4 addresses are priced as shown in Amazon VPC pricing page. This means that the subnets must have an associated route table with an internet gateway attached. For information about how to create and attach an internet gateway, see Enable VPC internet access using internet gateways in the Amazon VPC User Guide.

upvoted 1 times

FlyingHawk 4 months, 1 week ago

Unauthenticated access control must be off and at least one of the following access-control methods must be on: SASL/IAM, SASL/SCRAM, mTLS. For information about how to update the access-control method of a cluster, see Update security settings of a Amazon MSK cluster.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: A

"You can turn on public access to an MSK cluster at no additional cost..."

To turn on public access to a cluster, first ensure that the cluster meets all of the following conditions:

- The subnets that are associated with the cluster must be public.
- Unauthenticated access control must be off and at least one of the following access-control methods must be on: SASL/IAM, SASL/SCRAM, mTLS.
- ..."

<https://docs.aws.amazon.com/msk/latest/developerguide/public-access.html>

upvoted 4 times

MatAlves 8 months, 3 weeks ago

There is no reference to a NEW VPC being required in the documentation. We can simply configure subnets in the existing VPC.

upvoted 2 times

Scheldon 11 months, 2 weeks ago

Selected Answer: A

AnswerA

I need to agree that answer will probably be Option A.

upvoted 1 times

Indrasis 1 year, 3 months ago

Selected Answer: A

A is correct

upvoted 2 times

Marunio 1 year, 3 months ago

Selected Answer: A

A, since Kafka is loadbalancing itself. - <https://dattell.com/data-architecture-blog/load-balancing-with->

kafka/#:~:text=Load%20balancing%20with%20Kafka%20is,partitions%20while%20preserving%20message%20ordering.

B - why create new VPC?

C / D - Kafka is loadbalacing itself, also NLB can't handle HTTPS.

upvoted 3 times

Andy_09 1 year, 4 months ago

Option A

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #717

Topic 1

A company wants to migrate an on-premises legacy application to AWS. The application ingests customer order files from an on-premises enterprise resource planning (ERP) system. The application then uploads the files to an SFTP server. The application uses a scheduled job that checks for order files every hour.

The company already has an AWS account that has connectivity to the on-premises network. The new application on AWS must support integration with the existing ERP system. The new application must be secure and resilient and must use the SFTP protocol to process orders from the ERP system immediately.

Which solution will meet these requirements?

- A. Create an AWS Transfer Family SFTP internet-facing server in two Availability Zones. Use Amazon S3 storage. Create an AWS Lambda function to process order files. Use S3 Event Notifications to send s3:ObjectCreated:* events to the Lambda function.
- B. Create an AWS Transfer Family SFTP internet-facing server in one Availability Zone. Use Amazon Elastic File System (Amazon EFS) storage. Create an AWS Lambda function to process order files. Use a Transfer Family managed workflow to invoke the Lambda function.
- C. Create an AWS Transfer Family SFTP internal server in two Availability Zones. Use Amazon Elastic File System (Amazon EFS) storage. Create an AWS Step Functions state machine to process order files. Use Amazon EventBridge Scheduler to invoke the state machine to periodically check Amazon EFS for order files.
- D. Create an AWS Transfer Family SFTP internal server in two Availability Zones. Use Amazon S3 storage. Create an AWS Lambda function to process order files. Use a Transfer Family managed workflow to invoke the Lambda function. **Most Voted**

Correct Answer: D

Community vote distribution

D (84%)

A (16%)

Comments

anikolov **Highly Voted** 1 year, 4 months ago

Selected Answer: D

D looks more secure over existing on-prem to AWS connection
-Transfer Family SFTP internal server in two Availability Zones

Transfer Family with internal server in two Availability Zones.

-Use Amazon S3 storage.

-Use a Transfer Family managed workflow to invoke the Lambda function"

upvoted 11 times

hajra313 1 year, 3 months ago

If the legacy application needs to ingest customer order files from an on-premises ERP system and upload them to an SFTP server, an internet-facing AWS Transfer Family SFTP server would be the appropriate choice.

In this scenario, the SFTP server needs to be accessible from the internet to facilitate the file transfer between the on-premises system and AWS. Therefore, an internet-facing server is required to securely receive the files.

upvoted 2 times

Mr_Marcus 1 year ago

"The company already has an AWS account that has connectivity to the on-premises network." Internal Server.

upvoted 2 times

bora4motion Most Recent 2 weeks, 6 days ago

Selected Answer: D

A is wrong, the company already has a connection between aws and on prem, this means a dx or s2s vpn, therefore you dont have to expose the sftp to internet and transfer the data over the internet.

upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

Selected Answer: D

A and D are similar, but A internet-facing server is less secure than D, also no need to use S3 Event Notifications to send s3:ObjectCreated:* events to the Lambda function because Transfer Family managed workflows can directly invoke Lambda..

<https://docs.aws.amazon.com/transfer/latest/userguide/transfer-workflows.html>

upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

AWS Transfer Family supports managed workflows for file processing. With managed workflows, you can kick off a workflow after a file has been transferred over SFTP, FTPS, or FTP. Using this feature, you can securely and cost effectively meet your compliance requirements for business-to-business (B2B) file exchanges by coordinating all the necessary steps required for file processing. In addition, you benefit from end-to-end auditing and visibility.

<https://aws.amazon.com/blogs/storage/customize-file-delivery-notifications-using-aws-transfer-family-managed-workflows/>

upvoted 1 times

ckhemani 6 months, 2 weeks ago

Selected Answer: D

A and B is out since it is public facing and we don't need public facing SFTP since Company already have Connection to AWS thro ON-PREMISES.

Initially i thought it will C since it was doing scheduled as they were doing scheduled job every hour on on-premises. Then the read thro question and saw they are saying with new solution, they want to process the file immediately after it is uploaded so D is correct since it will trigger the Lambda file as soon as file appears on S3.

upvoted 3 times

Scheldon 11 months, 2 weeks ago

Selected Answer: D

Answer D

upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: D

"order files from an on-premises enterprise resource planning (ERP)" - Therefore Internal Endpoint is enough, no need for Internet-facing, although Internet-facing also handles on-prem connections as well, but "most secure". Even tho we are talking about SecureFTP.... Very bad wording of the question... :(

Definitely S3 against EFS, so D should be the answer...

upvoted 3 times

sandordini 1 year, 1 month ago

Also: With managed workflows, you can kick off a workflow after a file has been transferred over SFTP
upvoted 2 times

Hung23 1 year, 2 months ago

Selected Answer: A

Correct answer is A because must support integration with existing erp system we need to choose sftp internal-facing
upvoted 3 times

buzzinmumbai 1 year, 2 months ago

Answer is D . Both A&D are right but the question says it must support integration with existing erp system. I believe you can use transfer family for the existing job onprem as well to check for files.
upvoted 1 times

alawada 1 year, 2 months ago

Selected Answer: D

has an AWS account that has connectivity to the on-premises network.
upvoted 2 times

xBUGx 1 year, 3 months ago

Selected Answer: D

The company already has an AWS account that has connectivity to the on-premises network. So no need internet.
upvoted 3 times

67a3f49 1 year, 3 months ago

I would go in D as it's internal network.
upvoted 2 times

NayeraB 1 year, 3 months ago

Selected Answer: A

I think A makes more sense
upvoted 2 times

Andy_09 1 year, 4 months ago

A is the correct option
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #718

Topic 1

A company's applications use Apache Hadoop and Apache Spark to process data on premises. The existing infrastructure is not scalable and is complex to manage.

A solutions architect must design a scalable solution that reduces operational complexity. The solution must keep the data processing on premises.

Which solution will meet these requirements?

A. Use AWS Site-to-Site VPN to access the on-premises Hadoop Distributed File System (HDFS) data and application. Use an Amazon EMR cluster to process the data.


B. Use AWS DataSync to connect to the on-premises Hadoop Distributed File System (HDFS) cluster. Create an Amazon EMR cluster to process the data.

C. Migrate the Apache Hadoop application and the Apache Spark application to Amazon EMR clusters on AWS Outposts. Use the EMR clusters to process the data. **Most Voted**

D. Use an AWS Snowball device to migrate the data to an Amazon S3 bucket. Create an Amazon EMR cluster to process the data.

Correct Answer: C

Community vote distribution



Comments

anikolov **Highly Voted** 1 year, 4 months ago

Selected Answer: C

C cover requirement: The solution must keep the data processing on premises
upvoted 14 times

Andy_09 **Highly Voted** 1 year, 4 months ago

I would go for option C, as data processing has to be done on premise.
upvoted 9 times

Scheldon **Most Recent** 11 months, 2 weeks ago

Selected Answer: C

Answer C

upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: C

Only solution to keep the processing on-prem.

upvoted 2 times

Hung23 1 year, 2 months ago

Selected Answer: B

Create an Amazon EMR Cluster: With the data now available in Amazon S3, the company can create an Amazon EMR cluster for data processing. EMR provides scalable Hadoop and Spark clusters that can process data stored in S3, enabling the company to leverage cloud-based processing resources while still keeping the data processing on premises.

upvoted 3 times

example_ 10 months, 3 weeks ago

Selected response : C

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #719

Topic 1

A company is migrating a large amount of data from on-premises storage to AWS. Windows, Mac, and Linux based Amazon EC2 instances in the same AWS Region will access the data by using SMB and NFS storage protocols. The company will access a portion of the data routinely. The company will access the remaining data infrequently.


The company needs to design a solution to host the data.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an Amazon Elastic File System (Amazon EFS) volume that uses EFS Intelligent-Tiering. Use AWS DataSync to migrate the data to the EFS volume.
- B. Create an Amazon FSx for ONTAP instance. Create an FSx for ONTAP file system with a root volume that uses the auto tiering policy. Migrate the data to the FSx for ONTAP volume. **Most Voted**
- C. Create an Amazon S3 bucket that uses S3 Intelligent-Tiering. Migrate the data to the S3 bucket by using an AWS Storage Gateway Amazon S3 File Gateway.
- D. Create an Amazon FSx for OpenZFS file system. Migrate the data to the new volume.

Correct Answer: B

Community vote distribution



Comments

ogerber **Highly Voted** 1 year, 3 months ago

Selected Answer: B

Amazon FsX for NetAPP ONTAP feature: Multi-protocol access to data using the Network File System (NFS), Server Message Block (SMB), and Internet Small Computer Systems Interface (iSCSI) protocols
upvoted 18 times

jaswantn **Highly Voted** 1 year, 3 months ago

option C SMB and NFS storage protocols -> S3 file gateway
upvoted 7 times

FlyingHawk 4 months, 2 weeks ago

Initially I thought C is correct, but Windows, Mac, and Linux based Amazon EC2 instances in the same AWS Region will access the data by using SMB and NFS storage protocols, S3 does not provide the smb and NFS access to these ec2 instances, so c is incorrect, B is correct.

upvoted 3 times

MatAlves 8 months, 3 weeks ago

"S3 File Gateway is used for on-premises data intensive applications that need file protocol access to objects in S3. "

<https://aws.amazon.com/storagegateway/file/s3/>

upvoted 2 times

MatAlves 8 months, 3 weeks ago

The company is MIGRATING data from on-premises to AWS.

Amazon FSx for NetApp ONTAP offers high-performance file storage that's broadly accessible from Linux, Windows, and macOS compute instances via the industry-standard NFS, SMB, iSCSI, and NVMe-over-TCP protocols.

<https://aws.amazon.com/fsx/netapp-ontap/features/>

upvoted 3 times

MatAlves Most Recent 8 months, 3 weeks ago

Selected Answer: B

Amazon FSx for NetApp ONTAP offers high-performance file storage that's broadly accessible from Linux, Windows, and macOS compute instances via the industry-standard NFS, SMB, iSCSI, and NVMe-over-TCP protocols.

<https://aws.amazon.com/fsx/netapp-ontap/features/>

"S3 File Gateway is used for on-premises data intensive applications that need file protocol access to objects in S3. "

<https://aws.amazon.com/storagegateway/file/s3/>

upvoted 4 times

Scheldon 11 months, 2 weeks ago

Selected Answer: B

Answer B

upvoted 2 times

rjjkc 1 year ago

Selected Answer: C

I think it is "C"

Not "B" - The question never indicates the company is using "NetApp ONTAP file systems", so I am not sure what it means by "Migrate the data to the FSx for ONTAP volume". Please correct if misunderstood.

"C" clearly indicates how to migrate the data to S3, the S3 Intelligent-Tiering addressed the access pattern in the question and you can SMB/NFS mount S3 bucket

<https://docs.aws.amazon.com/filegateway/latest/files3/using-smb-fileshare.html>

<https://docs.aws.amazon.com/filegateway/latest/files3/GettingStartedAccessFileShare.html>

upvoted 1 times

TwinSpark 1 year ago

Selected Answer: B

FSx for ONTAP support NFS and SMB Protocol, even AWS Storage Gateway Amazon S3 File Gateway support them but it is used to connect on premises devices to file on s3, not to connect ec2 instances in the same aws region

upvoted 3 times

Linuslin 1 year ago

Selected Answer: B

Amazon FSx for NetApp ONTAP provides fully managed shared storage in the AWS Cloud with the popular data access and management capabilities of ONTAP.

Move workloads running on NetApp or other NFS/SMB/iSCSI servers to AWS without modifying application code or how you manage data.

And FsX for NetAPP ONTAP support "Reducing storage costs with automatic and intelligent storage tiering."

<https://aws.amazon.com/tw/fsx/netapp-ontap/faqs/#product-faqs#netapp-ontap-faq#reducing-storage-costs-with-automatic-and-intelligent-storage-tiering>

upvoted 4 times

camps 1 year, 2 months ago

it's C

upvoted 1 times

TruthWS 1 year, 2 months ago

B - FSx for ONTAP support SMB and NFS

upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: B

Amazon FSx for NetAPP ONTAP feature: Multi-protocol access to data using the Network File System (NFS), Server Message Block (SMB), and Internet Small Computer Systems Interface (iSCSI) protocols

Option C: make no sense I see it as a distractor

upvoted 2 times

Kezuko 1 year, 2 months ago

Selected Answer: C

Both B and C works, but it seems like C has a least operational overhead

upvoted 2 times

rondellidell 1 year, 2 months ago

The company will access the remaining data infrequently."

upvoted 1 times

Kezuko 1 year, 2 months ago

<https://www.amazonaws.cn/en/storagegateway/faqs/#:~:text=The%20Amazon%20S3%20File%20Gateway,be%20directly%20accessed%20in%20S3.>

upvoted 1 times

dkw2342 1 year, 2 months ago

It's B, option C makes no sense.

1. "Migrate the data to the S3 bucket using an AWS Storage Gateway Amazon S3 File Gateway." -> Nothing about running the gateway to access the files via SMB and NFS afterwards.

2. Even if you ignore this, the S3 File Gateway requires a virtual appliance to be deployed (on EC2 in this case), which contradicts the "LEAST operational overhead" requirement.

upvoted 2 times

Indrasis 1 year, 3 months ago

Selected Answer: C

Option C looks correct.

"The company will access a portion of the data routinely. The company will access the remaining data infrequently."

upvoted 3 times

Appon 1 year, 3 months ago

Selected Answer: B

option B

upvoted 2 times

MattBJ 1 year, 3 months ago

Selected Answer: C

C is correct

upvoted 3 times

hajra313 1 year, 3 months ago

Option A and D do not support SMB and NFS file system . Option b looks correvt

upvoted 5 times

Andy_09 1 year, 4 months ago

Option with S3 usage looks corrcet

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #720

Topic 1

A manufacturing company runs its report generation application on AWS. The application generates each report in about 20 minutes. The application is built as a monolith that runs on a single Amazon EC2 instance. The application requires frequent updates to its tightly coupled modules. The application becomes complex to maintain as the company adds new features.

Each time the company patches a software module, the application experiences downtime. Report generation must restart from the beginning after any interruptions. The company wants to redesign the application so that the application can be flexible, scalable, and gradually improved. The company wants to minimize application downtime.

Which solution will meet these requirements?

- A. Run the application on AWS Lambda as a single function with maximum provisioned concurrency.
- B. Run the application on Amazon EC2 Spot Instances as microservices with a Spot Fleet default allocation strategy.
- C. Run the application on Amazon Elastic Container Service (Amazon ECS) as microservices with service auto scaling.
Most Voted
- D. Run the application on AWS Elastic Beanstalk as a single application environment with an all-at-once deployment strategy.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Andy_09 **Highly Voted** 1 year, 4 months ago

Microservices using ECS
upvoted 10 times

MatAlves **Most Recent** 8 months, 3 weeks ago

Selected Answer: C

Monolith -> Microservices = ECS.
upvoted 1 times

Hung23 1 year, 2 months ago

Selected Answer: C

Sure 100%

upvoted 3 times

asdfcdsxdcf 1 year, 3 months ago

Selected Answer: C

Microservices using Elastic Container Service is correct

upvoted 1 times

Indrasis 1 year, 3 months ago

Selected Answer: C

C is correct

upvoted 1 times

Typewriter101 1 year, 3 months ago

Selected Answer: C

B will not help

spot instances provide cost savings but using it for a stateful task isn't right cause spot instances can be interrupted

upvoted 1 times

sandordini 1 year, 1 month ago

Correct answer but incorrect reasoning. Spot fleet includes on-demand AND spot instances to provide the desired capacity.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #721

Topic 1

A company wants to rearchitect a large-scale web application to a serverless microservices architecture. The application uses Amazon EC2 instances and is written in Python.


The company selected one component of the web application to test as a microservice. The component supports hundreds of requests each second. The company wants to create and test the microservice on an AWS solution that supports Python. The solution must also scale automatically and require minimal infrastructure and minimal operational support.

Which solution will meet these requirements?

- A. Use a Spot Fleet with auto scaling of EC2 instances that run the most recent Amazon Linux operating system.
- B. Use an AWS Elastic Beanstalk web server environment that has high availability configured.
- C. Use Amazon Elastic Kubernetes Service (Amazon EKS). Launch Auto Scaling groups of self-managed EC2 instances.
- D. Use an AWS Lambda function that runs custom developed code. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

Andy_09 **Highly Voted** 1 year, 4 months ago

Lambda looks like a better option
upvoted 11 times

Typewriter101 **Highly Voted** 1 year, 3 months ago

Selected Answer: D

Lambda
serverless, scalable, minimal infrastructure, handling hundreds of requests per second
upvoted 8 times

sandordini **Most Recent** 1 year, 1 month ago

Selected Answer: D

A: auto-scaling of EC2 instances = Lot of overhead + Infra

A: auto-scaling of EC2 instances - lot of overhead + infra

B: The company selected one component of the web application to test as a microservice. The component supports hundreds of requests each second. > lastic Beanstalk is a bad choice if you need worker processes. The whole point of a worker process is to perform a task in the background without slowing down your main web app. But Elastic Beanstalk doesn't support this option in a scalable way.

Also, they want to test just 1 selected microservice and I think it's a bit of overkill to do it using Elastic Beanstalk. Happy to be challenged though!

C: self-managed EC2 instances > infra + operational overhead

D: Lambda supports Python, microservice should be quicker than 15 mins, worst case scenario the test will fail.. (that's the purpose tests are conducted for anyway..)

I'd go for D

upvoted 5 times

gsgdga 1 year, 2 months ago

Selected Answer: C

microservice => EKS, ECS

upvoted 1 times

LuongTo 7 months, 2 weeks ago

containerized application then go with EKS, ECS is absolutely yes, but there would be more solution for microservices e.g. lambda

upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: C

C is the correct answer. The best way to deploy microservice is to use container-based service

upvoted 1 times

MatAlves 8 months, 3 weeks ago

"Maintain nodes yourself with self-managed nodes"

<https://docs.aws.amazon.com/eks/latest/userguide/worker.html>

upvoted 1 times

dkw2342 1 year, 2 months ago

Microservices doesn't automatically mean ECS or EKS. Read the question again: "Serverless" clearly contradicts "self-managed EC2 instances".

D is the only option that fits the criteria.

upvoted 2 times

rubiteb 1 year, 3 months ago

Best answer is C.

The application is a large-scale we app as mentioned in the question.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

You cannot have 'minimal infrastructure and minimal operational support' with 'Auto Scaling groups of self-managed EC2 instances.'

upvoted 2 times

rubiteb 1 year, 3 months ago

I mean B for Elastic Beans Stalk not C. EBS is the best solution for running large-scale application.

upvoted 1 times

Umuntu 1 year, 4 months ago

C is the correct answer. The best way to deploy microservice is to use container-based service such as EKS or ECS. So C is great

upvoted 3 times

Typewriter101 1 year, 3 months ago

Using ECS or EKS involves managing cluster and ec2 which will increase the infrastructure and operational overhead compared to lambda which is serverless.
upvoted 2 times

Andy_09 1 year, 4 months ago

EBS for minimal infra maintenance
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #722

Topic 1

A company has an AWS Direct Connect connection from its on-premises location to an AWS account. The AWS account has 30 different VPCs in the same AWS Region. The VPCs use private virtual interfaces (VIFs). Each VPC has a CIDR block that does not overlap with other networks under the company's control.

The company wants to centrally manage the networking architecture while still allowing each VPC to communicate with all other VPCs and on-premises networks.

Which solution will meet these requirements with the LEAST amount of operational overhead?

- A. Create a transit gateway, and associate the Direct Connect connection with a new transit VIF. Turn on the transit gateway's route propagation feature. **Most Voted**
- B. Create a Direct Connect gateway. Recreate the private VIFs to use the new gateway. Associate each VPC by creating new virtual private gateways.
- C. Create a transit VPCConnect the Direct Connect connection to the transit VPCCreate a peering connection between all other VPCs in the Region. Update the route tables.
- D. Create AWS Site-to-Site VPN connections from on premises to each VPC. Ensure that both VPN tunnels are UP for each connection. Turn on the route propagation feature.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Andy_09 **Highly Voted** 1 year, 4 months ago

Option A
upvoted 7 times

Umntu **Highly Voted** 1 year, 4 months ago

A is the best solution
upvoted 6 times

Besisco **Most Recent** 3 months, 4 weeks ago

Selected Answer: A

A, but... why there is not mentioned transit gateway attachments?

if no attachments, how VPCs are going to communicate onprem through the transit gateway?

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: A

"You can use AWS Direct Connect gateway to connect your Direct Connect connection over a transit virtual interface to the VPCs or VPNs that are attached to your transit gateway. You associate a Direct Connect gateway with the transit gateway. Then, create a transit virtual interface for your AWS Direct Connect connection to the Direct Connect gateway."

<https://docs.aws.amazon.com/directconnect/latest/UserGuide/direct-connect-transit-gateways.html>

upvoted 3 times

alawada 1 year, 2 months ago

Selected Answer: A

Turn on the transit gateway's route propagation feature.

upvoted 2 times

cedser8 1 year, 3 months ago

<https://docs.aws.amazon.com/directconnect/latest/UserGuide/direct-connect-gateways-intro.html>

upvoted 2 times

Typewriter101 1 year, 3 months ago

Selected Answer: A

transit gateway -> hub and spoke

upvoted 6 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #723

Topic 1

A company has applications that run on Amazon EC2 instances. The EC2 instances connect to Amazon RDS databases by using an IAM role that has associated policies. The company wants to use AWS Systems Manager to patch the EC2 instances without disrupting the running applications.

Which solution will meet these requirements?

- A. Create a new IAM role. Attach the AmazonSSMManagedInstanceCore policy to the new IAM role. Attach the new IAM role to the EC2 instances and the existing IAM role.
- B. Create an IAM user. Attach the AmazonSSMManagedInstanceCore policy to the IAM user. Configure Systems Manager to use the IAM user to manage the EC2 instances.
- C. Enable Default Host Configuration Management in Systems Manager to manage the EC2 instances. **Most Voted**
- D. Remove the existing policies from the existing IAM role. Add the AmazonSSMManagedInstanceCore policy to the existing IAM role.

Correct Answer: C*Community vote distribution*

C (80%)

A (20%)

Comments**jaswantn** **Highly Voted** 1 year, 3 months ago

option C....Default Host Management Configuration creates and applies a default IAM role to ensure that Systems Manager has permissions to manage all instances in the Region and perform automated patch scans using Patch Manager.

upvoted 13 times

Pics00094 **Highly Voted** 1 year, 3 months ago**Selected Answer: C**

C is the answer

upvoted 6 times

FlyingHawk **Most Recent** 5 months, 1 week ago**Selected Answer: C**

<https://docs.aws.amazon.com/systems-manager/latest/userguide/setup-instance-permissions.html> DHMC eliminates the need

<https://docs.aws.amazon.com/systems-manager/latest/userguide/setup-instance-permissions.html>, which eliminates the need to manually create and attach a new IAM role for Systems Manager, reducing operational overhead.

upvoted 2 times

LeonSauveterre 5 months, 3 weeks ago

Selected Answer: C

Well.. Actually I don't know what to choose. A or C ??

A - EC2 instances can have only one IAM role attached at a time. Whether attaching multiple roles is possible or not depends on whether AWS allows such configuration in your environment.

B - Systems Manager manages EC2 instances through the instance's attached IAM role, not an IAM user.

C - It doesn't address the core requirement of maintaining the existing RDS connectivity. It still requires the EC2 instances to have the necessary IAM permissions to be managed by SSM. The IAM role associated with the EC2 instance must still have the AmazonSSMManagedInstanceCore policy attached to enable Systems Manager capabilities.

D - Removing the existing policies would break the application's ability to connect to the RDS database. This directly contradicts the requirement of not disrupting running applications.

upvoted 1 times

FlyingHawk 5 months, 1 week ago

A is the way before having C, details see this blog:

<https://aws.amazon.com/blogs/mt/enable-management-of-your-amazon-ec2-instances-in-aws-systems-manager-using-default-host-management-configuration/>

upvoted 2 times

rosanna 5 months, 3 weeks ago

Selected Answer: A

I'd vote for A because the request here is not to disrupt the existing workload, meaning the existing IAM role must be intact with the addition to the new permission set that gives SSM patch manager patching capabilities.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: C

"The Default Host Management Configuration setting allows AWS Systems Manager to manage your Amazon EC2 instances automatically as managed instances.

Default Host Management Configuration makes it possible to manage EC2 instances without your having to manually create an AWS Identity and Access Management (IAM) instance profile. Instead, Default Host Management Configuration creates and applies a default IAM role to ensure that Systems Manager has permissions to manage all instances in the AWS account and AWS Region where it's activated."

upvoted 3 times

MatAlves 8 months, 3 weeks ago

<https://docs.aws.amazon.com/systems-manager/latest/userguide/fleet-manager-default-host-management-configuration.html>

upvoted 2 times

88f8032 1 year, 1 month ago

Selected Answer: A

i think A

upvoted 2 times

NayeraB 1 year, 3 months ago

So is C same as A, but automated?

upvoted 2 times

LeonSauveterre 5 months, 3 weeks ago

No, A is impossible because EC2 instances can have only one IAM role attached at a time.

upvoted 1 times

osmk 1 year, 3 months ago

C is fine
upvoted 2 times

Andy_09 1 year, 4 months ago

C is a better option
upvoted 3 times

Andy_09 1 year, 4 months ago

Correct answer A
upvoted 3 times

arunkpskpm 1 year, 3 months ago

"Attach the new IAM role to the EC2 instances and the existing IAM role" - You can't attach multiple policies to an EC2 instance. So A is wrong.
upvoted 6 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #724

Topic 1

A company runs container applications by using Amazon Elastic Kubernetes Service (Amazon EKS) and the Kubernetes Horizontal Pod Autoscaler. The workload is not consistent throughout the day. A solutions architect notices that the number of nodes does not automatically scale out when the existing nodes have reached maximum capacity in the cluster, which causes performance issues.

Which solution will resolve this issue with the LEAST administrative overhead?

- A. Scale out the nodes by tracking the memory usage.
- B. Use the Kubernetes Cluster Autoscaler to manage the number of nodes in the cluster. **Most Voted**
- C. Use an AWS Lambda function to resize the EKS cluster automatically.
- D. Use an Amazon EC2 Auto Scaling group to distribute the workload.

Correct Answer: B

Community vote distribution

B (100%)

Comments

FlyingHawk 5 months, 1 week ago

I think the better solution now is using Karpenter or turn on EKS Auto mode
<https://docs.aws.amazon.com/eks/latest/userguide/autoscaling.html>
 upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: B

Refer to

<https://www.examtopycs.com/discussions/amazon/view/109702-exam-aws-certified-solutions-architect-associate-saa-c03/>
 upvoted 2 times

Scheldon 11 months, 2 weeks ago

Selected Answer: B

AnswerB

Using of Kubernetes Cluster Autoscaler seems to be the best solution here

upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: B

When the workload increases and existing nodes reach maximum capacity, the Cluster Autoscaler detects the need for additional nodes and requests them from the underlying AWS infrastructure.

upvoted 2 times

osmk 1 year, 3 months ago

Selected Answer: B

Bcorrect

upvoted 2 times

Naveena_Devanga 1 year, 3 months ago

B is the correct answer. The Kubernetes Cluster Autoscaler automatically adjusts the number of nodes in your cluster when pods fail or are rescheduled onto other nodes. The Cluster Autoscaler uses Auto Scaling groups

upvoted 4 times

jaswantn 1 year, 3 months ago

option B.

upvoted 2 times

Andy_09 1 year, 4 months ago

Kubernetes Cluster Autoscaler looks correct

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #725

Topic 1

A company maintains about 300 TB in Amazon S3 Standard storage month after month. The S3 objects are each typically around 50 GB in size and are frequently replaced with multipart uploads by their global application. The number and size of S3 objects remain constant, but the company's S3 storage costs are increasing each month.

How should a solutions architect reduce costs in this situation?

- A. Switch from multipart uploads to Amazon S3 Transfer Acceleration.
- B. Enable an S3 Lifecycle policy that deletes incomplete multipart uploads. **Most Voted**
- C. Configure S3 inventory to prevent objects from being archived too quickly.
- D. Configure Amazon CloudFront to reduce the number of objects stored in Amazon S3.

Correct Answer: B*Community vote distribution*B (100%)**Comments****MatAlves** 8 months, 3 weeks ago**Selected Answer: B**

B - No brainer.
upvoted 1 times

MatAlves 8 months, 3 weeks ago

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/mpu-abort-incomplete-mpu-lifecycle-config.html>
upvoted 1 times

alawada 1 year, 2 months ago**Selected Answer: B**

Optimize multipart uploads to reduce costs associated with storing incomplete multipart upload parts. Ensure that multipart uploads are completed and the parts are assembled into complete objects in a timely manner to avoid unnecessary storage costs.
upvoted 4 times

Tuneawriter101 1 year, 3 months ago

typewriter101 1 year, 8 months ago

Selected Answer: B

when primary concern is cost and the data transfer multipart upload may be the more cost-effective than S3 transfer acceleration. So switching to s3 TA is won't be reasonable.

upvoted 4 times

Umuntu 1 year, 4 months ago

Option B is correct

upvoted 2 times

Andy_09 1 year, 4 months ago

Option B

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #726


Topic 1

A company has deployed a multiplayer game for mobile devices. The game requires live location tracking of players based on latitude and longitude. The data store for the game must support rapid updates and retrieval of locations.

The game uses an Amazon RDS for PostgreSQL DB instance with read replicas to store the location data. During peak usage periods, the database is unable to maintain the performance that is needed for reading and writing updates. The game's user base is increasing rapidly.

What should a solutions architect do to improve the performance of the data tier?

- A. Take a snapshot of the existing DB instance. Restore the snapshot with Multi-AZ enabled.
- B. Migrate from Amazon RDS to Amazon OpenSearch Service with OpenSearch Dashboards.
- C. Deploy Amazon DynamoDB Accelerator (DAX) in front of the existing DB instance. Modify the game to use DAX.
- D. Deploy an Amazon ElastiCache for Redis cluster in front of the existing DB instance. Modify the game to use Redis.

Most Voted**Correct Answer: D***Community vote distribution***Comments****Rhydian25** **Highly Voted** 11 months, 2 weeks ago**Selected Answer: D**

Why are people voting for C? PostgreSQL is a relational DB. DynamoDB is NoSQL.

It makes no sense
upvoted 7 times

1e22522 10 months ago

its bezos with his alt accounts.
upvoted 4 times

Dantecito **Most Recent** 3 months, 3 weeks ago

Selected Answer: D

To add more to all the answer Redis has built-in support for geospatial data, making it ideal for updates and retrieval of locations.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: D

Refer to <https://www.examttopics.com/discussions/amazon/view/53854-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Specific to Redis, ElastiCache lets you "scale in" or "scale out" both reads and writes. Cluster mode offers added shard support, enabling write scaling.

<https://aws.amazon.com/blogs/database/building-a-real-time-gaming-leaderboard-with-amazon-elasticache-for-redis/>

upvoted 2 times

Lin878 11 months, 2 weeks ago

Selected Answer: D

I confuse, Is DAX working with RDS?

upvoted 1 times

Jacky_S 11 months, 2 weeks ago

Selected Answer: D

The answer C is not making any sense with "Deploy Amazon DynamoDB Accelerator (DAX) in front of the existing DB instance.", because AWS DynamoDB is a DBaaS.

upvoted 2 times

Scheldon 11 months, 2 weeks ago

Selected Answer: D

AnswerD

ElastiCache is a fully managed, in-memory caching service that provides microsecond read and write latencies that support flexible, real-time use cases.

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/creating-elasticache-cluster-with-RDS-settings.html>

upvoted 2 times

Scheldon 11 months, 2 weeks ago

Beside DAX is for DynamoDB and I think it will not work with RDS

upvoted 1 times

ug56c 11 months, 3 weeks ago

Selected Answer: D

Amazon ElastiCache for Redis for RDS

upvoted 1 times

sheilawu 1 year ago

Selected Answer: C

" writing updates" so it should be DAX.

upvoted 2 times

f07ed8f 1 year ago

Selected Answer: C

Vote for C (DAX) as ElastiCache for Redis cluster only helps on read operation but not write.

upvoted 4 times

f07ed8f 1 year ago

Vote for C (DAX) as ElastiCache for Redis cluster only helps on read operation but not write operation.
upvoted 1 times

sirajtr47 1 year, 1 month ago

Performance > > Amazon ElastiCache for Redis cluster
upvoted 4 times

Rylz 1 year ago

Yeah, but this is an online game that need to update everything instantly so i don't think that caching is the right method here think about it, you are making movments in the online game its like loading new data all the time, so each time you move you load the cache for next time usage. what if you dont need to use it again?
i dont think elasticache is the right method here so it leave us with DAX
what do you think?
upvoted 1 times

Jacky_S 11 months, 2 weeks ago

i agree with DAX, but how it can be deploying in front of DB instance? Since that is a DBaaS.
upvoted 1 times

FZA24 1 year, 3 months ago

Selected Answer: D

D looks correct
upvoted 2 times

Umntu 1 year, 4 months ago

D looks correct
upvoted 4 times

Andy_09 1 year, 4 months ago

Looks correct
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #727

Topic 1

A company stores critical data in Amazon DynamoDB tables in the company's AWS account. An IT administrator accidentally deleted a DynamoDB table. The deletion caused a significant loss of data and disrupted the company's operations. The company wants to prevent this type of disruption in the future.

Which solution will meet this requirement with the LEAST operational overhead?

- A. Configure a trail in AWS CloudTrail. Create an Amazon EventBridge rule for delete actions. Create an AWS Lambda function to automatically restore deleted DynamoDB tables.
- B. Create a backup and restore plan for the DynamoDB tables. Recover the DynamoDB tables manually.
- C. Configure deletion protection on the DynamoDB tables. **Most Voted**
- D. Enable point-in-time recovery on the DynamoDB tables.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**BillaRanga** **Highly Voted** 9 months, 3 weeks ago**Selected Answer: C**

<https://aws.amazon.com/about-aws/whats-new/2023/03/amazon-dynamodb-table-deletion-protection/>

Deletion protection is now available for Amazon DynamoDB tables in all AWS Regions. DynamoDB now makes it possible for you to protect your tables from accidental deletion when performing regular table management operations. When creating new tables or managing existing tables, authorized administrators can set the deletion protection property for each table, which will govern whether a table can be deleted.

upvoted 15 times

BillaRanga 9 months, 3 weeks ago

Option B and D talks about recovering but not preventing. A is toooooo much work

upvoted 3 times

Andy_09 **Highly Voted** 10 months, 1 week ago

Option C

Option C
upvoted 5 times

Typewriter101 **Most Recent** 9 months, 3 weeks ago

Selected Answer: C

B involves more operations.
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #728

Topic 1

A company has an on-premises data center that is running out of storage capacity. The company wants to migrate its storage infrastructure to AWS while minimizing bandwidth costs. The solution must allow for immediate retrieval of data at no additional cost.

How can these requirements be met?

- A. Deploy Amazon S3 Glacier Vault and enable expedited retrieval. Enable provisioned retrieval capacity for the workload.
- B. Deploy AWS Storage Gateway using cached volumes. Use Storage Gateway to store data in Amazon S3 while retaining copies of frequently accessed data subsets locally. **Most Voted**
- C. Deploy AWS Storage Gateway using stored volumes to store data locally. Use Storage Gateway to asynchronously back up point-in-time snapshots of the data to Amazon S3.
- D. Deploy AWS Direct Connect to connect with the on-premises data center. Configure AWS Storage Gateway to store data locally. Use Storage Gateway to asynchronously back up point-in-time snapshots of the data to Amazon S3.

Correct Answer: B*Community vote distribution***Comments****67a3f49** **Highly Voted** 1 year, 3 months ago

B is the correct one because:

"A company has an on-premises data center that is running out of storage capacity".

So when they keep data on-premis and do the backup to S3 they'll run out of data and this is not their purpose.

upvoted 10 times

Sergantus 7 months ago

Storage Gateway in Stored mode DOS does NOT improve capacity – the main copy of data is stored on the gateway. It's B for that same reason

upvoted 1 times

JCVDB23 **Highly Voted** 1 year, 2 months ago**Selected Answer: B**

B. Deploy AWS Storage Gateway using cached volumes. Use Storage Gateway to store data in Amazon S3 while retaining copies of frequently accessed data subsets locally.

AWS Storage Gateway's cached volumes let you use Amazon S3 as your primary data storage while retaining frequently accessed data locally in your storage gateway. Cached volumes minimize the need to scale your on-premises storage infrastructure, while still providing your applications with low-latency access to their frequently accessed data. All data transferred between your gateway and AWS storage is encrypted for security. You can also save on data transfer costs as AWS Storage Gateway compresses all data transferred between the gateway and AWS, allowing you to store more data in AWS while reducing your data transfer costs.

upvoted 7 times

TEC65 Most Recent 1 month, 1 week ago

Selected Answer: B

they want to limit onsite storage

upvoted 1 times

tch 4 months, 2 weeks ago

Selected Answer: B

In the cached Volume Gateway mode, your primary data is stored in Amazon S3, while retaining your frequently accessed data locally in the cache for low latency access. In the stored Volume Gateway mode, your primary data is stored locally and your entire dataset is available for low latency access on premises while also asynchronously getting backed up to Amazon S3

upvoted 2 times

LeonSauveterre 5 months, 3 weeks ago

Selected Answer: B

Another headache one. Well, I'll go for B for now.

A - Really cost-effective, but immediate retrievals incurs additional costs.

B - Less optimal for datasets where most data is accessed frequently and must be immediately available locally.

C - The on-prem data center "is running out of storage capacity". Then why store everything locally? Why even migrate storage infrastructure to AWS...

D - Like C, and also, bandwidth costs are not acceptable.

upvoted 1 times

Lin878 11 months, 2 weeks ago

Selected Answer: C

I vote "C" because Question doesn't mention to access frequent data. If they want access frequent data, I will vote "B" with cached volume.

upvoted 1 times

1e22522 10 months ago

you don't fool me bezos boy

upvoted 3 times

Scheldon 11 months, 2 weeks ago

Selected Answer: B

AnswerB

I think B answer is the best option here. We will store only data which are frequently accessed and all other we will sent to the cloud. So we will have access to all data but hence most frequently accessed data will be stored in On-Premises Cache we will not pay a lot of additionally \$\$ for data transfer if any.

upvoted 1 times

ike001 11 months, 4 weeks ago

C is the correct. answer

upvoted 1 times

mohammadthainat 1 year, 2 months ago

Selected Answer: B

1- "The company wants to migrate its storage infrastructure to AWS" -> B as data will be migrated to AWS.

2- "The solution must allow for immediate retrieval of data at no additional cost." ->> B as data will be stored in S3 Standard storage class which provide immediate data retrieval

upvoted 3 times

gsgdga 1 year, 2 months ago

Selected Answer: C

immediate retrieval of data → should have full data set on-premises => stored volumes AWS Storage Gateway

upvoted 4 times

rubiteb 1 year, 3 months ago

B - as the company is migrating their data to AWS so data has to be stored in the cloud.

upvoted 2 times

osmk 1 year, 3 months ago

C>>>

Cached Mode: In this mode, your primary data resides in Amazon S3, while frequently accessed data is cached locally for low-latency access.

Stored Mode: Here, your entire dataset is stored locally, allowing low-latency access on premises. Simultaneously, the data is asynchronously backed up to Amazon S3.

upvoted 4 times

MatAlves 8 months, 3 weeks ago

How does that address the fact the company is "running out of storage capacity"?

upvoted 1 times

BillaRanga 1 year, 3 months ago

Selected Answer: C

D -> It takes One month to set up AWS Direct Connect setup

A -> No sense as it talks nothing about On-Prem

B -> Cached volume only stores frequently access data On-Prem, But requirement tells "Data" so we assume it tells All data

C -> Correct, as Stored volumes stores everything in Storage Gateway On-Prem while asynchronously backing up to the cloud

upvoted 5 times

sandordini 1 year, 1 month ago

D: It never said one month would be a problem.. Question doesn't state a matter of urgency, but it still stores the data on-prem, and synchronizes to AWS.

C: The same issue as D. Stores data locally, but our on-prem storage is full. That's why the company wants cloud.

A: Has retrieval costs.

upvoted 1 times

xBUGx 1 year, 2 months ago

i was voting for C, but C doesn't solve on-prem out of capacity issue.

upvoted 1 times

jaswantn 1 year, 3 months ago

option C... data being accessible through stored volume reduces bandwidth cost and provides immediate retrieval of data.

upvoted 2 times

Andy_09 1 year, 4 months ago

Option C, as it makes all the data available for low-latency access.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #729

Topic 1

A company runs a three-tier web application in a VPC across multiple Availability Zones. Amazon EC2 instances run in an Auto Scaling group for the application tier.

The company needs to make an automated scaling plan that will analyze each resource's daily and weekly historical workload trends. The configuration must scale resources appropriately according to both the forecast and live changes in utilization.

Which scaling strategy should a solutions architect recommend to meet these requirements?

- A. Implement dynamic scaling with step scaling based on average CPU utilization from the EC2 instances.
- B. Enable predictive scaling to forecast and scale. Configure dynamic scaling with target tracking **Most Voted**
- C. Create an automated scheduled scaling action based on the traffic patterns of the web application.
- D. Set up a simple scaling policy. Increase the cooldown period based on the EC2 instance startup time.

Correct Answer: B*Community vote distribution*B (100%)**Comments****sandordini** 7 months, 2 weeks ago**Selected Answer: B**

Not A: Only handles Dynamic scaling, not pattern-based/predictive scaling.

B: Both Predictive and dynamic

Not C: Manual version of predictive, lacks live circumstances..

Not D: The question doesn't talk about cool down period...

upvoted 3 times

alawada 8 months, 2 weeks ago

<https://aws.amazon.com/blogs/aws/new-predictive-scaling-for-ec2-powered-by-machine-learning/>

upvoted 2 times

BillaRanga 9 months, 3 weeks ago**Selected Answer: B**

By configuring dynamic scaling with target tracking, the company can automatically adjust resources based on the forecasted demand while also responding to live changes in utilization

upvoted 4 times

Andy_09 10 months, 1 week ago

Option B

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #730

Topic 1

A package delivery company has an application that uses Amazon EC2 instances and an Amazon Aurora MySQL DB cluster. As the application becomes more popular, EC2 instance usage increases only slightly. DB cluster usage increases at a much faster rate.


The company adds a read replica, which reduces the DB cluster usage for a short period of time. However, the load continues to increase. The operations that cause the increase in DB cluster usage are all repeated read statements that are related to delivery details. The company needs to alleviate the effect of repeated reads on the DB cluster.

Which solution will meet these requirements MOST cost-effectively?

- A. Implement an Amazon ElastiCache for Redis cluster between the application and the DB cluster. **Most Voted**
- B. Add an additional read replica to the DB cluster.
- C. Configure Aurora Auto Scaling for the Aurora read replicas.
- D. Modify the DB cluster to have multiple writer instances.

Correct Answer: A

Community vote distribution



Comments

Andy_09 **Highly Voted** 1 year, 4 months ago

Option A
upvoted 7 times

Scheldon **Most Recent** 11 months, 1 week ago

Selected Answer: C

AnswerC

After topic reconsideration I will go with Aurora autoscaling

To meet your connectivity and workload requirements, Aurora Auto Scaling dynamically adjusts the number of Aurora Replicas (reader DB instances) provisioned for an Aurora DB cluster. Aurora Auto Scaling is available for both Aurora MySQL and Aurora PostgreSQL. Aurora Auto Scaling enables your Aurora DB cluster to handle sudden increases in connectivity or workload. When

PostgreSQL. Aurora Auto Scaling enables your Aurora DB cluster to handle sudden increases in connectivity or workload. When the connectivity or workload decreases, Aurora Auto Scaling removes unnecessary Aurora Replicas so that you don't pay for unused provisioned DB instances.

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.Integrating.AutoScaling.html>

upvoted 1 times

MatAlves 8 months, 3 weeks ago

repeated reads = perfect scenario for CACHING.

How adding more reader instances will "alleviate the effect of repeated reads on the DB cluster"?

upvoted 2 times

Scheldon 11 months, 2 weeks ago

Selected Answer: A

AnswerA

upvoted 3 times

sandordini 1 year, 1 month ago

Selected Answer: A

A. Although Redis is not typically cheap, the question statement clearly shouts for a cached solution, which is Redis... Also, that's the only long-term solution, as we don't know anything about the volumes, scale of trends, etc...

upvoted 4 times

Kezuko 1 year, 2 months ago

Selected Answer: A

"repeated read statements" -> Cache layer

upvoted 3 times

BillaRanga 1 year, 3 months ago

Selected Answer: A

The question says, "The operations that cause the increase in DB cluster usage are all ****repeated read statements**** that are related to delivery details." - Read statements mean we can cache the results - hence, we need No read-replicas; we need only a cache layer to improve the performance.. Also, Adding read replicas costs money. The requirement is to meet them **MOST** cost-effectively

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #731

Topic 1

A company has an application that uses an Amazon DynamoDB table for storage. A solutions architect discovers that many requests to the table are not returning the latest data. The company's users have not reported any other issues with database performance. Latency is in an acceptable range.

Which design change should the solutions architect recommend?

- A. Add read replicas to the table.
- B. Use a global secondary index (GSI).
- C. Request strongly consistent reads for the table. **Most Voted**
- D. Request eventually consistent reads for the table.

Correct Answer: C

Community vote distribution



Comments

alawada **Highly Voted** 8 months, 2 weeks ago

Selected Answer: C

DynamoDB by default provides eventual consistency for read operations, which means that a query may not reflect the most recent data changes immediately after an update. Instead, it may take some time for the data to propagate across all replicas in the DynamoDB global table.

To ensure that read operations return the latest data and address the issue of stale data being returned to users, the solutions architect should recommend switching the read consistency level from eventually consistent reads to strongly consistent reads.

upvoted 7 times

BillaRanga **Most Recent** 9 months, 3 weeks ago

Selected Answer: C

Both tables and LSIs provide two read consistency options: eventually consistent (default) and strongly consistent reads.

1) Eventually Consistent Reads

Eventually consistent is the default read consistent model for all read operations. When issuing eventually consistent reads to a DynamoDB table or an index, the responses may not reflect the results of a recently completed write operation. If you repeat your read request after a short time, the response should eventually return the more recent item.

upvoted 4 times

BillaRanga 9 months, 3 weeks ago

2) Strongly Consistent Reads

Read operations such as GetItem, Query, and Scan provide an optional ConsistentRead parameter. If you set ConsistentRead to true, DynamoDB returns a response with the most up-to-date data, reflecting the updates from all prior write operations that were successful.

Hence it is C

A) Read-replicas are Async again, Which will persist the same problem.

B) Indexing will further cause latency, this has nothing to do with the question

upvoted 5 times

Andy_09 10 months, 1 week ago

Option C

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #732

Topic 1

A company has deployed its application on Amazon EC2 instances with an Amazon RDS database. The company used the principle of least privilege to configure the database access credentials. The company's security team wants to protect the application and the database from SQL injection and other web-based attacks.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use security groups and network ACLs to secure the database and application servers.
- B. Use AWS WAF to protect the application. Use RDS parameter groups to configure the security settings. **Most Voted**
- C. Use AWS Network Firewall to protect the application and the database.
- D. Use different database accounts in the application code for different functions. Avoid granting excessive privileges to the database users.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**Andy_09** **Highly Voted** 10 months, 1 week ago

Option B
upvoted 8 times

BillaRanga **Highly Voted** 9 months, 3 weeks ago**Selected Answer: B**

protect the application and the database from SQL injection and other web-based attacks. -> WAF
upvoted 5 times

03beafc **Most Recent** 7 months, 2 weeks ago

It's probably still B, but waf can't be attached directly to ec2's
upvoted 2 times

Typewriter101 9 months, 3 weeks ago**Selected Answer: B**

SQL injection -> WAF
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #733

Topic 1

An ecommerce company runs applications in AWS accounts that are part of an organization in AWS Organizations. The applications run on Amazon Aurora PostgreSQL databases across all the accounts. The company needs to prevent malicious activity and must identify abnormal failed and incomplete login attempts to the databases.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. Attach service control policies (SCPs) to the root of the organization to identify the failed login attempts.
- B. Enable the Amazon RDS Protection feature in Amazon GuardDuty for the member accounts of the organization.
Most Voted
- C. Publish the Aurora general logs to a log group in Amazon CloudWatch Logs. Export the log data to a central Amazon S3 bucket.
- D. Publish all the Aurora PostgreSQL database events in AWS CloudTrail to a central Amazon S3 bucket.

Correct Answer: B

Community vote distribution

B (100%)

Comments

BillaRanga **Highly Voted** 1 year, 3 months ago

Selected Answer: B

A -> SCPs are not for monitoring or logging

B-> correct

After you enable the RDS Protection feature, GuardDuty immediately starts monitoring RDS login activity from Aurora databases in your account. GuardDuty continuously monitors and profiles RDS login activity for suspicious activity, for example, unauthorized access to Aurora database in your account, from a previously unseen external actor.

upvoted 8 times

zinabu **Highly Voted** 1 year, 2 months ago

malicious activity=gurd duty

upvoted 5 times

Scheldon **Most Recent** 11 months, 2 weeks ago

Selected Answer: B

AnswerB

https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/CHAP_BestPractices.Security.html

Amazon GuardDuty is a threat detection service that helps protect your accounts, containers, workloads, and the data within your AWS environment. Using machine learning (ML) models, and anomaly and threat detection capabilities, GuardDuty continuously monitors different log sources and runtime activity to identify and prioritize potential security risks and malicious activities in your environment.

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/guard-duty-rds-protection.html>

upvoted 2 times

Naveena_Devanga 1 year, 3 months ago

B is the correct answer.

Amazon GuardDuty is a threat detection service that continuously monitors for malicious activity and unauthorized behavior to protect your Amazon Web Services accounts, workloads, and data stored in Amazon S3.

upvoted 4 times

Andy_09 1 year, 4 months ago

Option B

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #734

Topic 1

A company has an AWS Direct Connect connection from its corporate data center to its VPC in the us-east-1 Region. The company recently acquired a corporation that has several VPCs and a Direct Connect connection between its on-premises data center and the eu-west-2 Region. The CIDR blocks for the VPCs of the company and the corporation do not overlap. The company requires connectivity between two Regions and the data centers. The company needs a solution that is scalable while reducing operational overhead.

What should a solutions architect do to meet these requirements?

- A. Set up inter-Region VPC peering between the VPC in us-east-1 and the VPCs in eu-west-2.
- B. Create private virtual interfaces from the Direct Connect connection in us-east-1 to the VPCs in eu-west-2.
- C. Establish VPN appliances in a fully meshed VPN network hosted by Amazon EC2. Use AWS VPN CloudHub to send and receive data between the data centers and each VPC.
- D. Connect the existing Direct Connect connection to a Direct Connect gateway. Route traffic from the virtual private gateways of the VPCs in each Region to the Direct Connect gateway. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

JA2018 6 months, 1 week ago

Selected Answer: D

Connect the existing Direct Connect connection to a Direct Connect gateway. Route traffic from the virtual private gateways of the VPCs in each Region to the Direct Connect gateway.

Explanation:

Scalability and Reduced Overhead:

Using a Direct Connect gateway allows for easy scaling by adding more virtual private gateways from different regions to the same Direct Connect gateway, eliminating the need to manage separate connections for each region and minimizing operational overhead.

Inter-Region Connectivity:

A Direct Connect gateway acts as a central hub for routing traffic between VPCs in different regions through the Direct Connect connection, providing the required connectivity between the US-East-1 and EU-West-2 regions

upvoted 1 times

BillaRanga 1 year, 3 months ago

Selected Answer: D

"If you want to set up a Direct Connect to one or more VPC in many different regions (same account), you must use a Direct Connect Gateway."

upvoted 4 times

BillaRanga 1 year, 3 months ago

CloudHub is a VPN (encrypted) connection, so it goes over the public Internet., Whereas DirectConnect is Private (but not encrypted). So CloudHub is not suited for this useCase

upvoted 3 times

jaswantn 1 year, 3 months ago

option D

upvoted 2 times

Andy_09 1 year, 4 months ago

Changing to Option D for simpler implementation.

upvoted 3 times

1e22522 10 months ago

sure "Andy"

upvoted 3 times

Andy_09 1 year, 4 months ago

Option A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #735

Topic 1

A company is developing a mobile game that streams score updates to a backend processor and then posts results on a leaderboard. A solutions architect needs to design a solution that can handle large traffic spikes, process the mobile game updates in order of receipt, and store the processed updates in a highly available database. The company also wants to minimize the management overhead required to maintain the solution.

What should the solutions architect do to meet these requirements?

- A. Push score updates to Amazon Kinesis Data Streams. Process the updates in Kinesis Data Streams with AWS Lambda. Store the processed updates in Amazon DynamoDB. **Most Voted**
- B. Push score updates to Amazon Kinesis Data Streams. Process the updates with a fleet of Amazon EC2 instances set up for Auto Scaling. Store the processed updates in Amazon Redshift.
- C. Push score updates to an Amazon Simple Notification Service (Amazon SNS) topic. Subscribe an AWS Lambda function to the SNS topic to process the updates. Store the processed updates in a SQL database running on Amazon EC2.
- D. Push score updates to an Amazon Simple Queue Service (Amazon SQS) queue. Use a fleet of Amazon EC2 instances with Auto Scaling to process the updates in the SQS queue. Store the processed updates in an Amazon RDS Multi-AZ DB instance.

Correct Answer: A

Community vote distribution

A (100%)

Comments

BillaRanga **Highly Voted** 1 year, 3 months ago

Selected Answer: A

requirement -1: "Stream + process in order + Minimum Overhead" = Kinesis Data Stream + Lambda
requirement-2: "Highly available database + Min Management overhead" = DynamoDb

Setting Up Ec2 instance or MultiAZ DB = overhead
upvoted 12 times

Andy_09 **Highly Voted** 1 year, 4 months ago

Option A