Option A
upvoted 8 times

Scheldon Most Recent 11 months, 2 weeks ago

Selected Answer: A

AnswerA
upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: A

Even though it looks like SQS, but EC2 and Multi-AZ DB fail when it comes to minimal operational overhead.
upvoted 4 times

mohammadthainat 1 year, 2 months ago

Selected Answer: A

easy one: mobile game ->> DynamoDB
upvoted 2 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #736 *Topic 1*

A company has multiple AWS accounts with applications deployed in the us-west-2 Region. Application logs are stored within Amazon S3 buckets in each account. The company wants to build a centralized log analysis solution that uses a single S3 bucket. Logs must not leave us-west-2, and the company wants to incur minimal operational overhead.

Which solution meets these requirements and is MOST cost-effective?

A. Create an S3 Lifecycle policy that copies the objects from one of the application S3 buckets to the centralized S3 bucket.

B. Use S3 Same-Region Replication to replicate logs from the S3 buckets to another S3 bucket in us-west-2. Use this S3 bucket for log analysis. **Most Voted**

C. Write a script that uses the PutObject API operation every day to copy the entire contents of the buckets to another S3 bucket in us-west-2. Use this S3 bucket for log analysis.

D. Write AWS Lambda functions in these accounts that are triggered every time logs are delivered to the S3 buckets (s3:ObjectCreated:* event). Copy the logs to another S3 bucket in us-west-2. Use this S3 bucket for log analysis.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

BillaRanga Highly Voted 1 year, 3 months ago

Selected Answer: B

The main Use case of S3 same region replication is "log aggregation, live replication between production and test accounts".

upvoted 12 times

Andy_09 Highly Voted 1 year, 4 months ago

Option B

upvoted 6 times

zdi561 Most Recent 3 months, 4 weeks ago

Selected Answer: B

A and B can copy, but B is cheaper, it is free.

upvoted 1 times

ckhemani 6 months, 2 weeks ago

Selected Answer: B

Option C and D is out ot question since it required operation overhead.
Between A and B, Lifecycle is used to transfer between S3storage classes and not within same S3 Buckets. So B is correct answer.

upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: B

Needs to be B

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #737 *Topic 1*

A company has an application that delivers on-demand training videos to students around the world. The application also allows authorized content developers to upload videos. The data is stored in an Amazon S3 bucket in the us-east-2 Region.

The company has created an S3 bucket in the eu-west-2 Region and an S3 bucket in the ap-southeast-1 Region. The company wants to replicate the data to the new S3 buckets. The company needs to minimize latency for developers who upload videos and students who stream videos near eu-west-2 and ap-southeast-1.

Which combination of steps will meet these requirements with the FEWEST changes to the application? (Choose two.)

A. Configure one-way replication from the us-east-2 S3 bucket to the eu-west-2 S3 bucket. Configure one-way replication from the us-east-2 S3 bucket to the ap-southeast-1 S3 bucket.

B. Configure one-way replication from the us-east-2 S3 bucket to the eu-west-2 S3 bucket. Configure one-way replication from the eu-west-2 S3 bucket to the ap-southeast-1 S3 bucket.

C. Configure two-way (bidirectional) replication among the S3 buckets that are in all three Regions. **Most Voted**

D. Create an S3 Multi-Region Access Point. Modify the application to use the Amazon Resource Name (ARN) of the Multi-Region Access Point for video streaming. Do not modify the application for video uploads.

E. Create an S3 Multi-Region Access Point. Modify the application to use the Amazon Resource Name (ARN) of the Multi-Region Access Point for video streaming and uploads. **Most Voted**

Correct Answer: *CE*

Community vote distribution

CE (47%) AD (41%) Other (12%)

Comments

BillaRanga Highly Voted 1 year, 3 months ago

Selected Answer: CE

To keep replication in SYNC across all three regions, we use Bi-directional.
Multi-Region Access Point for video streaming and uploads. -> uploads to nearest Low latency region and Bi-directional replication will keep other two regions in SYNC this reducing the upload and streaming latency
upvoted 14 times

bujuman 1 year, 1 month ago

For confirmation purposes: https://aws.amazon.com/s3/features/multi-region-access-points/

upvoted 3 times

Dz1990 Most Recent 2 weeks, 2 days ago

Selected Answer: CE

□ **Grammar Breakdown:**
```
"minimize latency for [developers who upload videos] and [students who stream videos] near eu-west-2 and ap-southeast-1"
```

The phrase "near eu-west-2 and ap-southeast-1" modifies **BOTH** developers and students!

□ **Corrected Answer: CE**

Why CE is Actually Correct:

Option C - Bidirectional Replication:
- □ Allows developers to upload to their nearest region (low latency)
- □ Content automatically syncs to all other regions
- □ No single point of failure for uploads

Option E - MRAP for Everything:
- □ **Fewest application changes**: Single endpoint for all operations
- □ Automatic routing for both uploads AND downloads
- □ Future-proof and scalable

upvoted 1 times

srcntpc 3 weeks, 1 day ago

Selected Answer: AD

A & D
A: Sets up replication to eu-west-2 and ap-southeast-1.
D: Uses MRAP to optimize streaming latency without modifying upload logic.

upvoted 1 times

sk1974 3 months, 1 week ago

Selected Answer: AD

I will go with A & D too since that option will require least amount of re-work . In reality , I would have preferred B for the copy option since Ap-Sotheast-1 is closer to eu-west , but , that could involve additional work

upvoted 1 times

robotgeek 4 months, 1 week ago

Selected Answer: A

Question does not make a lot of sense with this responses
1) If the goal is to modify the app as little as possible, options D and E don't present much of a difference. If you modify the app for uploads, it will cost you the same to modify it for streaming (even more so with an MRAP, since you only need to change the global upload endpoint, which is a trivial change).
2) It is not suggested that the need is "to upload from anywhere in the world" but rather to copy to other regions for streaming (download), which an MRAP does not do per-se.
3) If the idea of the question is to create an MRAP (D or E) and then set up replication rules in the MRAP (C), it doesn't make sense to replicate from US-West or from Asia-Pacific, but only the other way around.
4) Option A stands on its own (but it's a multiple-choice answer), and even though the initial replication latency is slow, uploading video tutorials is not critical.
5) Option B involves unnecessary latency.

upvoted 1 times

robotgeek 4 months, 1 week ago

So the idea is you put an MRAP and then replicate... so basically the "FEWEST changes to the application" is either to confuse or poorly constructed

upvoted 1 times

FlyingHawk 5 months, 1 week ago

Selected Answer: AD

The data is stored in an Amazon S3 bucket in the us-east-2 Region implies the content developer is in us-east-2 region. If the content developers are in the us-east-2 Region, then the solution is A and D.
upvoted 3 times

kajiyatta 5 months, 2 weeks ago

Selected Answer: AD

This is an example of the kind of joke problem AWS provides. The information needed to derive justification is vague and it is not clear which region the uploading developer will be using.
Time spent on these problems is wasted.
upvoted 4 times

LeonSauveterre 5 months, 3 weeks ago

Selected Answer: AD

I agree with @kbgsgsgs and here's why:

What the scenario is about: Currenly, videos are in us, so they're uploaded to us. Then, we want to replicate them to eu and ap. People near eu and ap should be fast delivered with the videos they need.

So with "FEWEST changes", all we need is replicate from us to both eu and ap, which is option B. Option C works but not fewest changes.

The uploads are directly to us, so option E actually induces more latency. You may argue that if we can upload videos directly to eu/ap, it would be faster. Yes, but it requires modifying the application for both streaming and upload workflows, so much more complicated than option D, and definitely not fewest changes.
upvoted 2 times

LeonSauveterre 5 months, 3 weeks ago

Sorry, typo. All we need is replicate from us to both eu and ap, which is option A, not option B. You can share you ideas here if you disagree
upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: CD

My personal take: C & D

C => Although stem did not specify which region S# bucket the developers will upload their contents, this will allow new contents in 1 S3 bucket to replicate over to the other 2 S3 buckets.

D => uploads needs to be performed to the first region only (Closest???) and accessed by remaining two S3 buckets.
upvoted 1 times

JA2018 6 months, 1 week ago

but I am confused by Google AI search's response:

The correct answers to this question are D. Create an S3 Multi-Region Access Point. Modify the application to use the Amazon Resource Name (ARN) of the Multi-Region Access Point for video streaming and E. Create an S3 Multi-Region Access Point. Modify the application to use the Amazon Resource Name (ARN) of the Multi-Region Access Point for video streaming and uploads.
upvoted 1 times

JA2018 6 months, 1 week ago

Explanation:
Minimizing changes to the application:
The question explicitly asks for the solution with the "fewest changes to the application.". Option D and E only require modifying the application to use the Multi-Region Access Point ARN, while options A, B, and C would involve additional replication configurations, making them less ideal.
Multi-Region Access Point benefits:
An S3 Multi-Region Access Point allows users to access data from any region through a single endpoint, effectively minimizing latency for users in different regions without requiring data replication across multiple buckets.
upvoted 1 times

JA2018 6 months, 1 week ago

Why not the other options:
Options A and B:
These involve one-way replication, which would still require users in some regions to access data from a geographically distant bucket, leading to higher latency.
Option C:
Bidirectional replication would create unnecessary data redundancy and additional network traffic, making it less efficient.
Key takeaway: When you need to minimize latency for users in different regions while minimizing application changes, an S3 Multi-Region Access Point is the best solution.
upvoted 1 times

JA2018 6 months, 1 week ago

is it just me or am I missing something in this question?
upvoted 1 times

kbgsgsgs 8 months, 2 weeks ago

Selected Answer: AD

They are simply trying to replicate to a new S3 bucket. I don't see why it needs to be bidirectional. Also, since the problem assumes that the content developer with permission is the one uploading, it seems like there needs to be a way to centralize the upload without modifying the application.
upvoted 3 times

1166ae3 11 months ago

Selected Answer: AE

Since developer upload video to us-east-2, by configuring one-way replication directly from us-east-2 to eu-west-2 and from us-east-2 to ap-southeast-1, you ensure that each region has the latest data without additional replication hops.
upvoted 1 times

Scheldon 11 months, 2 weeks ago

Selected Answer: CE

AnswerCE

From my understanding Video uploads can happen near new regions hence to speed up that operation we need to upload to nearest region, hence I would choose option E, and for the same reason we need to be able to replicate data from any of new region to old one and oposite, hence we we need bidirectional (two-way) replication
upvoted 1 times

lenotc 1 year, 2 months ago

Selected Answer: CD

FEWEST changes to the application
D -> MRAP can upload the appropriate S3 bucket
C -> two-way -> to worry about anything
obs: I believe this question dubious, amphibological
upvoted 1 times

67a3f49 1 year, 3 months ago

There is no information where the upload should be performed. If files will be uploaded to first region then:

AD because:
A -> content uploaded to the primary bucket in us-east-2 is automatically replicated to the other regions, minimizing latency for users accessing content near those regions.
D -> uploads needs to be performed to the first region only and accessed to remaining two

Otherwise CE
upvoted 4 times

Salilgen 5 months, 1 week ago

You are right but I think files are only uploaded to first region.
The question states "A company has an application that delivers on-demand training videos to students around the world". It doesn't state anything about developer.
Furthermore, in this case both AD and CE work but AD requires fewest changes to the application
upvoted 2 times

Andy_09 1 year, 4 months ago

Correct answer CE

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #738 *Topic 1*

A company has a new mobile app. Anywhere in the world, users can see local news on topics they choose. Users also can post photos and videos from inside the app.

Users access content often in the first minutes after the content is posted. New content quickly replaces older content, and then the older content disappears. The local nature of the news means that users consume 90% of the content within the AWS Region where it is uploaded.

Which solution will optimize the user experience by providing the LOWEST latency for content uploads?

A. Upload and store content in Amazon S3. Use Amazon CloudFront for the uploads.

B. Upload and store content in Amazon S3. Use S3 Transfer Acceleration for the uploads. **Most Voted**

C. Upload content to Amazon EC2 instances in the Region that is closest to the user. Copy the data to Amazon S3.

D. Upload and store content in Amazon S3 in the Region that is closest to the user. Use multiple distributions of Amazon CloudFront.

Correct Answer: *B*

Community vote distribution

B (88%) D (12%)

Comments

Cali182 Highly Voted 1 year, 4 months ago

Selected Answer: B

Cloudfront is for reading not for uploading
Option B

upvoted 17 times

Salilgen 5 months, 1 week ago

I don't think so
https://aws.amazon.com/blogs/aws/amazon-cloudfront-content-uploads-post-put-other-methods/

upvoted 1 times

BillaRanga Highly Voted 1 year, 3 months ago

Selected Answer: B

Question says - " LOWEST latency for content uploads"
Hence Use S3 Transfer Acceleration for the uploads.
upvoted 13 times

flaviobrf Most Recent 10 months, 2 weeks ago

Selected Answer: D

Cloudfront can also upload data, not just for caching content
upvoted 2 times

flaviobrf 10 months, 2 weeks ago

Sorry, for this scenario, i believe that B its the correct
upvoted 1 times

ike001 11 months, 4 weeks ago

B is the answer
upvoted 2 times

[Removed] 1 year ago

option D. Upload and store content in Amazon S3 in the Region that is closest to the user. Use multiple distributions of Amazon CloudFront.
This solution ensures low-latency uploads by storing content in the nearest S3 region and provides fast access to users by distributing content through CloudFront edge locations.
upvoted 7 times

7fb06b3 1 year ago

Selected Answer: D

CloudFront does support upload acceleration
https://aws.amazon.com/blogs/aws/amazon-cloudfront-content-uploads-post-put-other-methods/
upvoted 3 times

TruthWS 1 year, 2 months ago

Option D
upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: B

Amazon S3 Transfer Acceleration utilizes Amazon CloudFront's globally distributed edge locations to accelerate content uploads to Amazon S3.
upvoted 3 times

xBUGx 1 year, 2 months ago

Selected Answer: B

S3TA is acturally using cloudfront's infrastructure.
So, yes B. Which is just an optimized solution for cloudfront itself.
upvoted 2 times

lpergorta 1 year, 2 months ago

Option D
Regional S3 Buckets: Storing content in S3 buckets located in the same Region as the user minimizes the physical distance the data needs to travel during upload, reducing latency.
CloudFront Distributions: CloudFront is a content delivery network (CDN) that caches content in edge locations around the world. By creating multiple CloudFront distributions with edge locations closest to users, the content can be served with minimal latency for downloads.
upvoted 3 times

Andy_09 1 year, 4 months ago

Option D

upvoted 4 times

jaswantn 1 year, 3 months ago

option B... S3 transfer acceleration for LOWEST latency for content uploads. question is not asking for low latency for content retrieval.
Happy to be corrected

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #739 *Topic 1*

A company is building a new application that uses serverless architecture. The architecture will consist of an Amazon API Gateway REST API and AWS Lambda functions to manage incoming requests.

The company wants to add a service that can send messages received from the API Gateway REST API to multiple target Lambda functions for processing. The service must offer message filtering that gives the target Lambda functions the ability to receive only the messages the functions need.

Which solution will meet these requirements with the LEAST operational overhead?

A. Send the requests from the API Gateway REST API to an Amazon Simple Notification Service (Amazon SNS) topic. Subscribe Amazon Simple Queue Service (Amazon SQS) queues to the SNS topic. Configure the target Lambda functions to poll the different SQS queues.

B. Send the requests from the API Gateway REST API to Amazon EventBridge. Configure EventBridge to invoke the target Lambda functions. **Most Voted**

C. Send the requests from the API Gateway REST API to Amazon Managed Streaming for Apache Kafka (Amazon MSK). Configure Amazon MSK to publish the messages to the target Lambda functions.

D. Send the requests from the API Gateway REST API to multiple Amazon Simple Queue Service (Amazon SQS) queues. Configure the target Lambda functions to poll the different SQS queues.

Correct Answer: *B*

Community vote distribution

B (49%) A (49%) D

Comments

Kezuko Highly Voted 1 year, 2 months ago

Selected Answer: A

"message filtering" = SNS

upvoted 14 times

AMEJack Highly Voted 6 months, 2 weeks ago

Selected Answer: B

The question is asking about one service. If you choose A, you need to build 3 services one per each SNS topic. Also, creating one Event Bridge with rules to forward to the target Lambda is easier than creating 3 topics with 3 SQS queues. The answer is definetly B.

upvoted 5 times

GOTJ Most Recent 2 months, 4 weeks ago

Selected Answer: A

I'm discarding option "B" for the following reasons

1. You can't send "requests" (messages) to EventBridge, at least not directly as option suggests. There must be an event involved in the process
2. Lack of configuration regarding to discrete Lambda invocations

Option "A", by its side, also lacks of the necessary configuration for discrete lambda invocation (use of subscriptor filter), but in my opinion is closer to what company needs:

https://docs.aws.amazon.com/sns/latest/dg/sns-message-filtering.html

So my vote goes to "A".

upvoted 1 times

Chen77 3 months, 3 weeks ago

Selected Answer: A

SNS supports fan-out: API Gateway sends messages to an SNS topic, which then forwards to multiple SQS queues.
Message filtering with SNS: SNS allows each SQS queue to receive only the relevant messages, reducing unnecessary processing.
SQS queues decouple workloads: Lambda functions pull messages only when ready, ensuring efficient and scalable processing.

upvoted 2 times

zdi561 3 months, 4 weeks ago

Selected Answer: B

A and B both can do message filtering but A only filters message by attributes and B can do filtering by message content too, and B is simpler

upvoted 1 times

hilker1983 5 months, 1 week ago

Selected Answer: B

LEAST operational overhead -> B better than A

upvoted 1 times

Denise123 5 months, 1 week ago

Selected Answer: B

The answer is B. On the other hand, the solution involving Amazon SNS, Amazon SQS, and Lambda polling (Option A) introduces more moving parts and operational overhead:
It requires managing and configuring multiple services (SNS, SQS, and Lambda).
Lambda functions need to continuously poll SQS queues, which can be less efficient and introduce additional latency.
Scaling and managing SQS queues and their associated resources (e.g., dead-letter queues) add operational complexity.
While the SNS-SQS-Lambda solution is still valid and widely used, the EventBridge-Lambda integration offers a more streamlined and serverless-native approach, aligning better with the company's serverless architecture goals and minimizing operational overhead.

upvoted 1 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: B

EventBridge applies filtering rules and invokes the appropriate Lambda functions directly, just like SNS + SQS. Option B is definitely "the LEAST operational overhead".

While there’s a 5-target-per-rule limit, you can work around it by creating additional rules. Each rule can directly invoke Lambda functions without the need for SQS queues or polling.

Let's not stick to option A just because it's the textbook exam-standard SNS+SQS combination. There are other options that will help you achieve the goal under different circumstances.

upvoted 2 times

upvoted 2 times

JA2018 6 months, 1 week ago

Selected Answer: B

agreed with AMEJack.... based on the stem, B required the LEAST operational overhead

upvoted 2 times

JA2018 6 months, 1 week ago

From Google AI Search:

The correct answer is B. Send the requests from the API Gateway REST API to Amazon EventBridge. Configure EventBridge to invoke the target Lambda functions.

upvoted 1 times

JA2018 6 months, 1 week ago

Explanation:
Least operational overhead:
EventBridge is designed specifically for event routing and filtering, making it the most efficient option for managing which Lambda functions receive messages based on specific criteria, minimizing operational overhead compared to the other choices.
Message filtering:
EventBridge allows you to define detailed rules with attributes and conditions to filter messages and send them only to the relevant Lambda functions, fulfilling the requirement for message filtering.
Scalability:
Both EventBridge and SNS can handle large volumes of messages, but EventBridge offers better management and control over message routing, especially when dealing with complex event patterns.

upvoted 1 times

JA2018 6 months, 1 week ago

Why the other options are less ideal:
A. SNS with SQS:
While SNS can be used for fan-out messaging, adding SQS queues introduces extra complexity and potential bottlenecks in managing multiple queues and subscriptions.
C. Amazon MSK:
While Kafka is a powerful messaging system, it might be overkill for this scenario due to its complexity and overhead for managing a dedicated Kafka cluster, especially if the message volume is not very high.
D. Multiple SQS queues:
This option would require additional logic to decide which SQS queue to send messages to based on filtering criteria, adding complexity to the architecture.
Key takeaway: When looking for a solution with the least operational overhead for message routing and filtering between an API Gateway and multiple Lambda functions, Amazon EventBridge is the most suitable choice.

upvoted 1 times

1166ae3 11 months ago

Selected Answer: B

LEAST operational overhead -> B better than A

upvoted 5 times

Scheldon 11 months, 2 weeks ago

Selected Answer: A

AnswerA

Hence EventBridge is a solution to handle events and we need to handle messages I believe option A is the best solution here

upvoted 3 times

3bdf1cc 12 months ago

https://aws.amazon.com/blogs/compute/capturing-client-events-using-amazon-api-gateway-and-amazon-eventbridge/

upvoted 2 times

BBR01 1 year, 1 month ago

Selected Answer: A

The main issue with B is that with Eventbrige, you can only define up to five targets for each rule

The main issue with B is that with Eventbrige, you can only define up to five targets for each rule.
https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-targets.html
upvoted 4 times

sandordini 1 year, 1 month ago

B: EventBridge reacts to events, not requests or messages.
C: I don't think so, but I don't know MSK well enough.

D: You can add a filter so that your function only processes Amazon SQS messages containing certain data parameters. but it will still receive, so I assume it's not what the question asks for.
Only A remains... But it still misses steps plus we are looking for the Least ops overhead..

I am confused..
upvoted 2 times

03beafc 1 year, 1 month ago

Selected Answer: B

Eventbridge + lambda is two services, sns + sqs + lambda is 3. Both can filter, but the config involved in eventbridge > lambda is easier
upvoted 3 times

MatAlves 8 months, 3 weeks ago

SNS provides built-in message filtering.
SNS -> SQS -> Lambda = very common Fanout Scenario
upvoted 2 times

AlvinC2024 1 year, 2 months ago

Selected Answer: D

Upload and store content in Amazon S3 in the Region that is closest to the user. Use multiple distributions of Amazon CloudFront. This approach ensures that uploads are quick, taking advantage of the geographical proximity of S3, while still leveraging CloudFront for efficient content delivery outside the local region if necessary. The local nature of the content consumption aligns with storing content in the closest region to the user, addressing the requirement that 90% of the content is consumed within the AWS Region where it is uploaded.
upvoted 1 times

TruthWS 1 year, 2 months ago

Option B - Eventbridge allow routing event from source to dest or multi dest you want
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #740 *Topic 1*

A company migrated millions of archival files to Amazon S3. A solutions architect needs to implement a solution that will encrypt all the archival data by using a customer-provided key. The solution must encrypt existing unencrypted objects and future objects.

Which solution will meet these requirements?

A. Create a list of unencrypted objects by filtering an Amazon S3 Inventory report. Configure an S3 Batch Operations job to encrypt the objects from the list with a server-side encryption with a customer-provided key (SSE-C). Configure the S3 default encryption feature to use a server-side encryption with a customer-provided key (SSE-C). **Most Voted**

B. Use S3 Storage Lens metrics to identify unencrypted S3 buckets. Configure the S3 default encryption feature to use a server-side encryption with AWS KMS keys (SSE-KMS).

C. Create a list of unencrypted objects by filtering the AWS usage report for Amazon S3. Configure an AWS Batch job to encrypt the objects from the list with a server-side encryption with AWS KMS keys (SSE-KMS). Configure the S3 default encryption feature to use a server-side encryption with AWS KMS keys (SSE-KMS).

D. Create a list of unencrypted objects by filtering the AWS usage report for Amazon S3. Configure the S3 default encryption feature to use a server-side encryption with a customer-provided key (SSE-C).

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

OX_HDR Highly Voted 1 year, 4 months ago

Selected Answer: A

A seems correct here.

https://aws.amazon.com/blogs/storage/encrypting-objects-with-amazon-s3-batch-operations/

upvoted 9 times

BillaRanga Highly Voted 1 year, 3 months ago

Selected Answer: A

S3 inventory list has "Encryption status" field so you can use this to filter the unencrypted objects. and use S3 batch to encrypt it with SSE-C key.

AWS Usage report does not provide details about encryption status of individual objects

upvoted 9 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: A

As stated, "by using a customer-provided key", which rules out B & C.

Option A uses 3 thingies:
1. S3 Inventory - Provides a native way to generate a list of unencrypted objects, making it scalable for millions of objects.
2. S3 Batch Operations - Encrypts existing objects efficiently.
3. S3 Default Encryption - Ensures future objects are automatically encrypted (with SSE-C).

upvoted 1 times

Scheldon 11 months, 2 weeks ago

Selected Answer: A

AnswerA

upvoted 2 times

ike001 11 months, 4 weeks ago

A is the answer

upvoted 2 times

jaswantn 1 year, 3 months ago

option B.... S3 Inventory report to check for unencrypted objects in s3 and then using Batch operation.

upvoted 1 times

mestule 1 year, 4 months ago

Selected Answer: A

The solution must encrypt existing unencrypted objects. Batch will do that.

upvoted 5 times

Andy_09 1 year, 4 months ago

Option B

upvoted 1 times

sk1974 3 months, 1 week ago

Need to use customer provided keys.So , B cannot be the answer

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #741 *Topic 1*

The DNS provider that hosts a company's domain name records is experiencing outages that cause service disruption for a website running on AWS. The company needs to migrate to a more resilient managed DNS service and wants the service to run on AWS.

What should a solutions architect do to rapidly migrate the DNS hosting service?

A. Create an Amazon Route 53 public hosted zone for the domain name. Import the zone file containing the domain records hosted by the previous provider. **Most Voted**

B. Create an Amazon Route 53 private hosted zone for the domain name. Import the zone file containing the domain records hosted by the previous provider.

C. Create a Simple AD directory in AWS. Enable zone transfer between the DNS provider and AWS Directory Service for Microsoft Active Directory for the domain records.

D. Create an Amazon Route 53 Resolver inbound endpoint in the VPC. Specify the IP addresses that the provider's DNS will forward DNS queries to. Configure the provider's DNS to forward DNS queries for the domain to the IP addresses that are specified in the inbound endpoint.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option A

upvoted 9 times

BillaRanga Highly Voted 1 year, 3 months ago

Selected Answer: A

A -> Correct as we need to route to a Company in public network.
B -> No, because it routes only within one or more VPC
C -> Added as a distractor
D -> Inbound resolver is for traffic from On-Prem to VPC

upvoted 7 times

aditianand 1 year ago

Hello Billaranga, I just bought this examtopics. My exam is on Jun 9th. Did we get questions from exam topics? How was the exam?

upvoted 3 times

ccceb01 Most Recent 9 months, 2 weeks ago

Selected Answer: A

A is the answer
https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/migrate-dns-domain-in-use.html

upvoted 2 times

pawanghujanamazon53 1 year, 1 month ago

Selected Answer: A

option A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #742 *Topic 1*

A company is building an application on AWS that connects to an Amazon RDS database. The company wants to manage the application configuration and to securely store and retrieve credentials for the database and other services.

Which solution will meet these requirements with the LEAST administrative overhead?

A. Use AWS AppConfig to store and manage the application configuration. Use AWS Secrets Manager to store and retrieve the credentials. **Most Voted**

B. Use AWS Lambda to store and manage the application configuration. Use AWS Systems Manager Parameter Store to store and retrieve the credentials.

C. Use an encrypted application configuration file. Store the file in Amazon S3 for the application configuration. Create another S3 file to store and retrieve the credentials.

D. Use AWS AppConfig to store and manage the application configuration. Use Amazon RDS to store and retrieve the credentials.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

BillaRanga Highly Voted 9 months, 3 weeks ago

Selected Answer: A

AppConfig useCase = You can use AWS AppConfig to deploy configuration data stored in the AWS AppConfig hosted configuration store, AWS Secrets Manager, Systems Manager Parameter Store, or Amazon S3.
So B and C are out.

use RDS to store credentials is not a good practise. So D is out.

Ans is A

upvoted 10 times

Andy_09 Highly Voted 10 months, 1 week ago

Option A

upvoted 8 times

Awsbeginner87 Most Recent 8 months, 1 week ago

Credentials= secrets Manager

upvoted 3 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #743 *Topic 1*

To meet security requirements, a company needs to encrypt all of its application data in transit while communicating with an Amazon RDS MySQL DB instance. A recent security audit revealed that encryption at rest is enabled using AWS Key Management Service (AWS KMS), but data in transit is not enabled.

What should a solutions architect do to satisfy the security requirements?

A. Enable IAM database authentication on the database.

B. Provide self-signed certificates. Use the certificates in all connections to the RDS instance.

C. Take a snapshot of the RDS instance. Restore the snapshot to a new instance with encryption enabled.

D. Download AWS-provided root certificates. Provide the certificates in all connections to the RDS instance. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (81%) A (19%)

Comments

BillaRanga Highly Voted 1 year, 3 months ago

Selected Answer: D

Amazon RDS creates an SSL certificate and installs the certificate on the DB instance when the instance is provisioned. So it is AWS provided.

upvoted 13 times

Andy_09 Highly Voted 1 year, 4 months ago

Option D

upvoted 5 times

Arad Most Recent 6 months, 1 week ago

Selected Answer: A

A is the correct answer

upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: D

Selected Answer: D

Why D is correct:
To encrypt data in transit to an RDS MySQL instance, you need to configure the database connection to use TLS/SSL, which requires using AWS-provided root certificates to establish a secure connection.

upvoted 3 times

JA2018 6 months, 1 week ago

Why other options are incorrect:
A (IAM database authentication): While important for user access control, this does not address data encryption in transit.
B (Self-signed certificates): Using self-signed certificates is not recommended for production environments as they cannot be verified by the client and might raise security concerns.
C (Snapshot and restore): Taking a snapshot and restoring to a new instance with encryption enabled only affects data at rest, not data in transit.

upvoted 2 times

JA2018 6 months, 1 week ago

From Google AI Search:

Key point: To encrypt data in transit to an RDS instance, ensure your application uses the appropriate AWS-provided root certificates to establish a secure TLS/SSL connection.

upvoted 1 times

Scheldon 11 months, 2 weeks ago

Selected Answer: D

AnswerD

upvoted 2 times

DAIYL 1 year, 1 month ago

Selected Answer: D

Even if IAM database authentication is enabled, clients still need to download and configure the AWS-provided root certificate to ensure a secure connection using SSL/TLS encryption. Without configuring the certificate, communication may not be fully encrypted, even with IAM authentication enabled.
https://docs.aws.amazon.com/zh_cn/AmazonRDS/latest/UserGuide/UsingWithRDS.SSL.html

upvoted 4 times

Nm55569 1 year ago

That's not in any of the answers - "Provide the certificates in all connections to the RDS instance." this doesn't make sense with option D - it's not saying configure to trust the CA. Answer can only be option A. Your link includes this "Optionally, your SSL/TLS connection can perform server identity verification by validating the server certificate installed on your database.". This you don't actually need to trust the using CA and can configure the app that way - the traffic is still encrypted though.

upvoted 3 times

Kezuko 1 year, 2 months ago

Selected Answer: A

A

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/UsingWithRDS.IAMDBAuth.html

upvoted 4 times

Sivaeas 1 year, 3 months ago

Optiona A:
IAM database authentication provides the following benefits:

Network traffic to and from the database is encrypted using Secure Socket Layer (SSL) or Transport Layer Security (TLS). For more information about using SSL/TLS with Amazon RDS, see Using SSL/TLS to encrypt a connection to a DB instance or cluster.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #744 *Topic 1*

A company is designing a new web service that will run on Amazon EC2 instances behind an Elastic Load Balancing (ELB) load balancer. However, many of the web service clients can only reach IP addresses authorized on their firewalls.

What should a solutions architect recommend to meet the clients' needs?

A. A Network Load Balancer with an associated Elastic IP address. **Most Voted**

B. An Application Load Balancer with an associated Elastic IP address.

C. An A record in an Amazon Route 53 hosted zone pointing to an Elastic IP address.

D. An EC2 instance with a public IP address running as a proxy in front of the load balancer.

Correct Answer: *A*

Community vote distribution

A (95%) C (5%)

Comments

67a3f49 Highly Voted 1 year, 3 months ago

A for sure. The same question was in "AWS Certified Solutions Architect Associate Practice Test 3" on Udemy. There was an explaination that NLB needs to be before ALB because only NLB can have static IP.

upvoted 16 times

BillaRanga Highly Voted 1 year, 3 months ago

Selected Answer: A

B -> Application Load Balancer cannot be assigned an Elastic IP address (static IP address).
C -> Its DNS after all, "Associated elastic IP" is what IP? Makes no sense
D -> "If you require a persistent public IP address that can be associated to and from instances as you require, use an Elastic IP address instead." PUBLIC IP of an EC2 is not persistent, although we can give an Elastic Ip, Using EC2 in front of a Load Balancer is tooooo much. What if it gets a million request? So to scale that EC2 you use another LB and an ASG>? This makes no sense

A is correct because a NLB can have an elastic IP and we can use this in our firewall as per the use case

upvoted 5 times

bora4motion Most Recent 2 weeks, 5 days ago

Selected Answer: A

Why not B? we're talking about http traffic not tcp / udp. Why not ALB instead of NLB?
upvoted 1 times

FlyingHawk 5 months, 1 week ago

Selected Answer: A

Both A and C are correct:

A Uses an NLB because it supports Elastic IP addresses, which are static IPs that can be whitelisted by clients.

C : Create an Alias record in Route 53 that points to the NLB.

An Alias record is a Route 53-specific feature that allows you to map a domain name (e.g., webservice.example.com) directly to the NLB without needing to use the NLB's IP address in the DNS record.

Clients can access the web service using the domain name (e.g., webservice.example.com), which resolves to the NLB's IP address.

However, in the context of exam, I will select A since the question explicitly mentions the use of an Elastic Load Balancer (ELB), and A directly addresses this by recommending a Network Load Balancer (NLB) with an Elastic IP address.
upvoted 2 times

Scheldon 11 months, 2 weeks ago

Selected Answer: A

AnswerA
upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: A

A - correct (Static ip can thereafter be used for client whitelisting)
Using a Network Load Balancer instead of a Classic Load Balancer has the following benefits:
Support for static IP addresses for the load balancer.
https://docs.aws.amazon.com/elasticloadbalancing/latest/network/introduction.html
upvoted 5 times

Sivaeas 1 year, 3 months ago

Selected Answer: A

Option A

Please look into the below for detailed explaination
https://www.scalefactory.com/blog/2021/12/13/aws-network-load-balancers-new-features/img/previously-firewall-egress.png
upvoted 3 times

PolarFox 1 year, 3 months ago

Selected Answer: C

Option C
upvoted 1 times

JA2018 6 months, 1 week ago

B. An Application Load Balancer with an associated Elastic IP address:
While an ALB can use an Elastic IP, it operates at the application layer and may not be suitable when the primary concern is client firewall restrictions requiring a specific IP.

C. An A record in an Amazon Route 53 hosted zone pointing to an Elastic IP address:
While Route 53 can manage DNS records, it does not directly address the need for a specific IP address accessible by clients with firewall restrictions.

D. An EC2 instance with a public IP address running as a proxy in front of the load balancer:
This adds unnecessary complexity and introduces a single point of failure, making it less than optimal.
upvoted 1 times

haira313 1 year, 3 months ago

Setting up an EC2 instance with a public IP address to act as a proxy in front of the load balancer allows clients with restricted IP access to connect to the web service. The EC2 instance can handle IP address whitelisting and proxy requests to the ELB load balancer, ensuring that only authorized clients can access the service. This solution provides flexibility and control over access while leveraging the scalability and availability benefits of ELB.

upvoted 1 times

BillaRanga 1 year, 3 months ago

Is this ChatGPT answer? Can you provide the AWS documentation link?

upvoted 3 times

Andy_09 1 year, 4 months ago

Option C

upvoted 2 times

jaswantn 1 year, 3 months ago

is there any valid justification for opting C? Glad to be informed, as these questions are tricky to answer.

upvoted 2 times

jaswantn 1 year, 3 months ago

My inclination is for Option D, but not 100 % sure

upvoted 1 times

JA2018 6 months, 1 week ago

D. An EC2 instance with a public IP address running as a proxy in front of the load balancer:
This adds unnecessary complexity and introduces a single point of failure, making it less than optimal.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #745 *Topic 1*

A company has established a new AWS account. The account is newly provisioned and no changes have been made to the default settings. The company is concerned about the security of the AWS account root user.

What should be done to secure the root user?

A. Create IAM users for daily administrative tasks. Disable the root user.

B. Create IAM users for daily administrative tasks. Enable multi-factor authentication on the root user. **Most Voted**

C. Generate an access key for the root user. Use the access key for daily administration tasks instead of the AWS Management Console.

D. Provide the root user credentials to the most senior solutions architect. Have the solutions architect use the root user for daily administration tasks.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

BillaRanga Highly Voted 9 months, 3 weeks ago

Selected Answer: B

"As a best practice, do not use the AWS account root user for any task where it's not required. Instead, create a new IAM user for each person that requires administrator access."
It's B :)

upvoted 12 times

Andy_09 Highly Voted 10 months, 1 week ago

Option B

upvoted 8 times

d401c0d Most Recent 7 months, 3 weeks ago

Selected Answer: B

D is just killing me. If we have reached this far, we all know it is Option B - "As a best practice, do not use the AWS account root user for any task where it's not required. Instead, create a new IAM user for each person that requires administrator access."

upvoted 2 times

Sivaeas 9 months ago

Selected Answer: B

its option B

upvoted 3 times

Naveena_Devanga 9 months, 3 weeks ago

Segregation of roles, also known as separation of duties (SoD), is a business control
that helps prevent security or privacy incidents and errors. Therefore, root access must never used for routine operational activities.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #746 *Topic 1*

A company is deploying an application that processes streaming data in near-real time. The company plans to use Amazon EC2 instances for the workload. The network architecture must be configurable to provide the lowest possible latency between nodes.

Which combination of network solutions will meet these requirements? (Choose two.)

A. Enable and configure enhanced networking on each EC2 instance. **Most Voted**

B. Group the EC2 instances in separate accounts.

C. Run the EC2 instances in a cluster placement group. **Most Voted**

D. Attach multiple elastic network interfaces to each EC2 instance.

E. Use Amazon Elastic Block Store (Amazon EBS) optimized instance types.

Correct Answer: *AC*

Community vote distribution

AC (100%)

Comments

mestule Highly Voted 10 months ago

Selected Answer: AC

A. Enable and configure enhanced networking on each EC2 instance. Enhanced networking provides higher bandwidth, higher packet per second (PPS) performance, and consistently lower inter-instance latencies.

C. Run the EC2 instances in a cluster placement group. A cluster placement group is a logical grouping of instances within a single Availability Zone. This configuration is recommended for applications that need low network latency, high network throughput, or both.

upvoted 12 times

Sivaeas Highly Voted 9 months ago

Selected Answer: AC

To reach speeds up to 10 Gbps between instances, launch your instances in a cluster placement group with the enhanced networking instance type. These instance types are placed physically close to each other. Instance types that are close to each other further reduces latency and improves transfer speeds.

upvoted 5 times

FlyingHawk Most Recent 5 months, 1 week ago

Selected Answer: AC

A: https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/enhanced-networking.html
C: https://aws.amazon.com/ec2/instance-types/#Instance_Features
Select EC2 instances support cluster networking when launched into a common cluster placement group. A cluster placement group provides low-latency networking between all instances in the cluster. The bandwidth an EC2 instance can utilize depends on the instance type and its networking performance specification.

upvoted 1 times

AmirBe 8 months, 3 weeks ago

AC
Use of Placement Groups: Utilize EC2 Placement Groups to ensure that instances are physically located close to each other within the same Availability Zone. This reduces the latency between instances by minimizing the distance data needs to travel.
Selection of EC2 Instance Types: Choose EC2 instance types optimized for low-latency networking, such as instances with enhanced networking capabilities like Elastic Network Adapter (ENA) or instances that support Amazon EC2 Nitro System. These instances provide high throughput and low latency networking performance.

upvoted 4 times

osmk 9 months, 3 weeks ago

what's AM?

upvoted 1 times

jaswantn 10 months ago

option C & E.
Option A is not viable as ..EC2 provides enhanced networking capabilities using single root I/O virtualization (SR-IOV) only on supported instance types.

upvoted 1 times

FlyingHawk 5 months, 1 week ago

Enhanced networking is available on most modern instance families, such as C5, M5, R5, P3, G4, and others.
Always check the AWS documentation https://aws.amazon.com/ec2/instance-types/ for the latest list of supported instance types:

upvoted 1 times

jaswantn 10 months ago

option E... EBS–optimized instance uses an optimized configuration

upvoted 1 times

sandordini 7 months, 2 weeks ago

C'mon, EBS is storage. The question does not deal with the storage solutions. Its a distractor...

upvoted 2 times

Andy_09 10 months, 1 week ago

Correct option should be CD

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #747 *Topic 1*

A financial services company wants to shut down two data centers and migrate more than 100 TB of data to AWS. The data has an intricate directory structure with millions of small files stored in deep hierarchies of subfolders. Most of the data is unstructured, and the company's file storage consists of SMB-based storage types from multiple vendors. The company does not want to change its applications to access the data after migration.

What should a solutions architect do to meet these requirements with the LEAST operational overhead?

A. Use AWS Direct Connect to migrate the data to Amazon S3.

B. Use AWS DataSync to migrate the data to Amazon FSx for Lustre.

C. Use AWS DataSync to migrate the data to Amazon FSx for Windows File Server. **Most Voted**

D. Use AWS Direct Connect to migrate the data on-premises file storage to an AWS Storage Gateway volume gateway.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

Sivaeas Highly Voted 1 year, 3 months ago

Selected Answer: C

AWS DataSync is a data transfer service that simplifies, automates, and accelerates moving and replicating data between on-premises storage systems and AWS storage services over the internet or AWS Direct Connect. DataSync can transfer your file data, and also file system metadata such as ownership, time stamps, and access permissions.

In DataSync, a location for Amazon FSx for Windows is an endpoint for an FSx for Windows File Server. You can transfer files between a location for Amazon FSx for Windows and a location for other file systems. For information, see Working with Locations in the AWS DataSync User Guide.

DataSync accesses your FSx for Windows File Server using the Server Message Block (SMB) protocol.

upvoted 6 times

aditianand 1 year ago

Does Fsx support SMB? I read prior posts that it only ONTAP supports SMB

upvoted 1 times

Andy_09 Highly Voted 1 year, 4 months ago

Option C

upvoted 6 times

Scheldon Most Recent 11 months, 2 weeks ago

Selected Answer: C

AnswerC

upvoted 2 times

Linuslin 1 year ago

Selected Answer: C

AWS Storage Gateway provides a standard set of storage protocols such as iSCSI, SMB, and NFS, which allow you to use AWS storage without rewriting your existing applications.--->A is not complete describe, so A is out.
https://aws.amazon.com/storagegateway/faqs/?nc=sn&loc=6

Only FSx for NetApp ONTAP and FSx for Windows File Server support SMB Protocol. --->B is out.
https://aws.amazon.com/tw/fsx/when-to-choose-fsx/

AWS Direct Connect is more expensive than AWS DataSync.--->D is out

C is the correct answer.

upvoted 2 times

Naveena_Devanga 1 year, 3 months ago

Correct Anwer is C
As most of the data is unstructured, and the company's file storage consists of SMB-based storage types from multiple vendors which is commonly a Windows-Linux file-sharing type so FSx for Windows File Server file systems completely meets the solution.

upvoted 3 times

ogerber 1 year, 3 months ago

Selected Answer: C

Option C since its SMB (windows) , and low operational effort so DataSync over Direct Connect

upvoted 4 times

osmk 1 year, 3 months ago

Selected Answer: C

https://docs.aws.amazon.com/datasync/latest/userguide/create-fsx-location.html

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #748 *Topic 1*

A company uses an organization in AWS Organizations to manage AWS accounts that contain applications. The company sets up a dedicated monitoring member account in the organization. The company wants to query and visualize observability data across the accounts by using Amazon CloudWatch.

Which solution will meet these requirements?

A. Enable CloudWatch cross-account observability for the monitoring account. Deploy an AWS CloudFormation template provided by the monitoring account in each AWS account to share the data with the monitoring account. Most Voted

B. Set up service control policies (SCPs) to provide access to CloudWatch in the monitoring account under the Organizations root organizational unit (OU).

C. Configure a new IAM user in the monitoring account. In each AWS account, configure an IAM policy to have access to query and visualize the CloudWatch data in the account. Attach the new IAM policy to the new IAM user.

D. Create a new IAM user in the monitoring account. Create cross-account IAM policies in each AWS account. Attach the IAM policies to the new IAM user.

Correct Answer: *A*

Community vote distribution

A (91%) C (9%)

Comments

jaswantn Highly Voted 1 year, 3 months ago

option A
below are the links to check both parts of option A.
https://docs.amazonaws.cn/en_us/AmazonCloudWatch/latest/monitoring/cloudwatch_crossaccount_dashboard.html

https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CloudWatch-Unified-Cross-Account-Setup.html#Unified-Cross-Account-SetupSource-SingleTemplate
upvoted 7 times

MatAlves 8 months, 3 weeks ago

"If you have multiple Amazon accounts, you can set up CloudWatch cross-account observability and then create rich cross-account dashboards in your monitoring accounts. You can seamlessly search, visualize, and analyze your metrics, logs, and traces without account boundaries."

traces without account boundaries.

Nice catch!
upvoted 3 times

Sivaeas Highly Voted 1 year, 3 months ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CloudWatch-Unified-Cross-Account.html
upvoted 6 times

ninasgx Most Recent 1 year, 3 months ago

Selected Answer: C

It's C
upvoted 1 times

osmk 1 year, 3 months ago

Selected Answer: A

https://docs.amazonaws.cn/en_us/AmazonCloudWatch/latest/monitoring/cloudwatch_crossaccount_dashboard.html
upvoted 4 times

Andy_09 1 year, 4 months ago

Option A
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #749 *Topic 1*

A company's website is used to sell products to the public. The site runs on Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer (ALB). There is also an Amazon CloudFront distribution, and AWS WAF is being used to protect against SQL injection attacks. The ALB is the origin for the CloudFront distribution. A recent review of security logs revealed an external malicious IP that needs to be blocked from accessing the website.

What should a solutions architect do to protect the application?

A. Modify the network ACL on the CloudFront distribution to add a deny rule for the malicious IP address.

B. Modify the configuration of AWS WAF to add an IP match condition to block the malicious IP address. **Most Voted**

C. Modify the network ACL for the EC2 instances in the target groups behind the ALB to deny the malicious IP address.

D. Modify the security groups for the EC2 instances in the target groups behind the ALB to deny the malicious IP address.

Correct Answer: *B*

Community vote distribution

B (87%) A (13%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option B

upvoted 12 times

mohammadthainat Highly Voted 1 year, 2 months ago

Selected Answer: B

in WAF you can define Web ACL (Web Access Control List) Rule:
IP Set: up to 10,000 IP addresses – use multiple Rules for more IPs

upvoted 5 times

FlyingHawk Most Recent 5 months, 1 week ago

Selected Answer: B

Network ACL is with VPC at subnet level, A and C are incorrect. security group for EC2 only supports allow rule, not deny rul, D is out. B is correct:
https://docs.aws.amazon.com/waf/latest/developerguide/waf-rule-statement-type-ipset-match.html

upvoted 2 times

Salilgen 5 months, 1 week ago

Selected Answer: B

https://aws.amazon.com/blogs/security/how-to-enhance-amazon-cloudfront-origin-security-with-aws-waf-and-aws-secrets-manager/

upvoted 1 times

bujuman 1 year, 1 month ago

Selected Answer: B

There was option from disttribution Security Tab ==> Request logs for the specified time range where someone could target an IP address and block it - which action won't do more than creating a block rule under the associated Web ACL- but function has vanished, i don't ask me why.
So the only feasable option in WEBACLv2 is to go for an Ipset and ad a WebACL ip match block condition.
I really liked the option A the first time i experimented it.

upvoted 4 times

xBUGx 1 year, 2 months ago

Selected Answer: A

You only need to block an IP. And Cloudfront is the first layer

upvoted 3 times

FlyingHawk 5 months, 1 week ago

CloudFront does not support network ACLs. Network ACLs are used with Amazon VPCs at the subnet level, not with CloudFront distributions.
https://docs.aws.amazon.com/vpc/latest/userguide/vpc-network-acls.html

upvoted 1 times

JA2018 6 months, 1 week ago

Modifying the CloudFront network ACL would block access at the CDN level. This could
impact legitimate users globally, not just the malicious IP.

upvoted 1 times

JA2018 6 months, 1 week ago

When dealing with web application security concerns, leverage the capabilities of AWS WAF to block malicious IPs rather than modifying network access controls at lower levels like EC2 security groups or network ACLs

upvoted 1 times

Sivaeas 1 year, 3 months ago

Selected Answer: B

The AWS WAF IP set match statement inspects the IP address of a web request against a set of IP addresses and address ranges. Use this to allow or block web requests based on the IP addresses that the requests originate from

upvoted 4 times

stephensimudemy 1 year, 3 months ago

Selected Answer: B

Option B

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #750 *Topic 1*

A company sets up an organization in AWS Organizations that contains 10 AWS accounts. A solutions architect must design a solution to provide access to the accounts for several thousand employees. The company has an existing identity provider (IdP). The company wants to use the existing IdP for authentication to AWS.

Which solution will meet these requirements?

A. Create IAM users for the employees in the required AWS accounts. Connect IAM users to the existing IdP. Configure federated authentication for the IAM users.

B. Set up AWS account root users with user email addresses and passwords that are synchronized from the existing IdP.

C. Configure AWS IAM Identity Center (AWS Single Sign-On). Connect IAM Identity Center to the existing IdP. Provision users and groups from the existing IdP. **Most Voted**

D. Use AWS Resource Access Manager (AWS RAM) to share access to the AWS accounts with the users in the existing IdP.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

osmk Highly Voted 9 months, 3 weeks ago

c--> Regardless of how you provision users, IAM Identity Center redirects the AWS Management Console, command line interface, and application authentication to your external IdP. IAM Identity Center then grants access to those resources based on policies you create in IAM Identity Center https://docs.aws.amazon.com/singlesignon/latest/userguide/manage-your-identity-source-idp.html#provisioning-when-external-idp

upvoted 9 times

ogerber Most Recent 9 months, 3 weeks ago

Selected Answer: C

Option C
https://docs.aws.amazon.com/singlesignon/latest/userguide/manage-your-identity-source-idp.html

upvoted 4 times

osmk 9 months, 3 weeks ago

Selected Answer: C

https://docs.aws.amazon.com/singlesignon/latest/userguide/manage-your-identity-source-idp.html#provisioning-when-external-idp

upvoted 3 times

Andy_09 10 months, 1 week ago

Option C

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #751 *Topic 1*

A solutions architect is designing an AWS Identity and Access Management (IAM) authorization model for a company's AWS account. The company has designated five specific employees to have full access to AWS services and resources in the AWS account.

The solutions architect has created an IAM user for each of the five designated employees and has created an IAM user group.

Which solution will meet these requirements?

A. Attach the AdministratorAccess resource-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group.

B. Attach the SystemAdministrator identity-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group.

C. Attach the AdministratorAccess identity-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group. **Most Voted**

D. Attach the SystemAdministrator resource-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

Scheldon Highly Voted 11 months, 2 weeks ago

Selected Answer: C

AnswerC

We need identity-based policy and if we will compare System Admin and Administrator Access policy it clear that SysAdmin have is allowing for limited amount of actions, where Admin Access simple allow for all actions.

https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_identity-vs-resource.html

https://docs.aws.amazon.com/aws-managed-policy/latest/reference/AdministratorAccess.html

https://docs.aws.amazon.com/aws-managed-policy/latest/reference/SystemAdministrator.html
upvoted 7 times

Linuslin Highly Voted 1 year ago

Selected Answer: C

The question says "full access to AWS services and resources in the AWS account" and "created an IAM user group."
You can see it is identity-based policy, not resource-based.--->A and D are out.
SystemAdministrator: Allow 28 of 412 services.--->B is out.
AdministratorAccess: Allow 412 of 412 services.--->C is the correct answer.

If you are curious about what a policy can allow for, just log in you AWS account and go to IAM-policies to find out.
upvoted 7 times

NSA_Poker Most Recent 11 months, 4 weeks ago

Selected Answer: C

(A & D) eliminated. Resource-based policies are attached to a resource NOT an IAM user, group, or role.
(B) eliminated. SystemAdministrator has fewer permissions than AdministratorAccess.
upvoted 5 times

MattBJ 1 year, 3 months ago

Selected Answer: C

C is the correct answer
upvoted 2 times

osmk 1 year, 3 months ago

Selected Answer: C

C>>>https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_manage-attach-detach.html
upvoted 3 times

osmk 1 year, 3 months ago

C>>>https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_manage-attach-detach.html
upvoted 3 times

Umuntu 1 year, 4 months ago

C looks correct
upvoted 3 times

Andy_09 1 year, 4 months ago

Option C
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #752 *Topic 1*

A company has a multi-tier payment processing application that is based on virtual machines (VMs). The communication between the tiers occurs asynchronously through a third-party middleware solution that guarantees exactly-once delivery.

The company needs a solution that requires the least amount of infrastructure management. The solution must guarantee exactly-once delivery for application messaging.

Which combination of actions will meet these requirements? (Choose two.)

A. Use AWS Lambda for the compute layers in the architecture. Most Voted

B. Use Amazon EC2 instances for the compute layers in the architecture.

C. Use Amazon Simple Notification Service (Amazon SNS) as the messaging component between the compute layers.

D. Use Amazon Simple Queue Service (Amazon SQS) FIFO queues as the messaging component between the compute layers. Most Voted

E. Use containers that are based on Amazon Elastic Kubernetes Service (Amazon EKS) for the compute layers in the architecture.

Correct Answer: *AD*

Community vote distribution

AD (100%)

Comments

osmk Highly Voted 1 year, 3 months ago

Selected Answer: AD

https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/FIFO-queues-exactly-once-processing.html
upvoted 6 times

Scheldon Highly Voted 11 months, 2 weeks ago

Selected Answer: AD

AnswerAD,

SQS FIFO will guarantee exactly one time execution for each operation. The problem is with processing as we do not know if

whole process will be closed in 15 min (TTL for Lambda).
I'm choosing Lambda as it is natural thing for Payment procesing in AWS but I'm not 100% sure
upvoted 5 times

tch Most Recent 5 months, 2 weeks ago

Selected Answer: AD

SQS is helpful for Asynchronous processing workflows.
upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: AD

"As its name suggests, exactly-once semantics means that each message is delivered precisely once. The message can neither be lost nor delivered twice (or more times)."

- SNS doesn't provide exactly-once delivery. Thus, we need SQS.
- To achieve "least amount of intra management", we go with Lambda for compute layer.

upvoted 2 times

Sivaeas 1 year, 3 months ago

Selected Answer: AD

Lamdba+SQS FIFO
upvoted 4 times

PolarFox 1 year, 3 months ago

someone please explain why the combination of D and E is not the correct?
upvoted 2 times

stephensimudemy 1 year, 3 months ago

because qn says 'least amount of infrastructure management'.
E is not.
upvoted 2 times

JA2018 6 months, 1 week ago

While EKS can provide scalability, it still requires managing container infrastructure. This adds complexity (more overhead in infra mgmt than A) to the solution
upvoted 1 times

jaswantn 1 year, 3 months ago

option A for payment processing.
option D for exactly once delivery.
upvoted 2 times

Umuntu 1 year, 4 months ago

CD IS THE BEST ANSWER
upvoted 1 times

hajra313 1 year, 4 months ago

a and d
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #753 | *Topic 1*

A company has a nightly batch processing routine that analyzes report files that an on-premises file system receives daily through SFTP. The company wants to move the solution to the AWS Cloud. The solution must be highly available and resilient. The solution also must minimize operational effort.

Which solution meets these requirements?

A. Deploy AWS Transfer for SFTP and an Amazon Elastic File System (Amazon EFS) file system for storage. Use an Amazon EC2 instance in an Auto Scaling group with a scheduled scaling policy to run the batch operation. **Most Voted**

B. Deploy an Amazon EC2 instance that runs Linux and an SFTP service. Use an Amazon Elastic Block Store (Amazon EBS) volume for storage. Use an Auto Scaling group with the minimum number of instances and desired number of instances set to 1.

C. Deploy an Amazon EC2 instance that runs Linux and an SFTP service. Use an Amazon Elastic File System (Amazon EFS) file system for storage. Use an Auto Scaling group with the minimum number of instances and desired number of instances set to 1.

D. Deploy AWS Transfer for SFTP and an Amazon S3 bucket for storage. Modify the application to pull the batch files from Amazon S3 to an Amazon EC2 instance for processing. Use an EC2 instance in an Auto Scaling group with a scheduled scaling policy to run the batch operation.

Correct Answer: *A*

Community vote distribution

A (76%) | D (21%) | C

Comments

Sivaeas Highly Voted 1 year, 3 months ago

Selected Answer: A

The Answer should be A not D because ...
Modify the application to pull the batch files from Amazon S3 to an Amazon EC2 instance for processing.--Why we need to do this when we can move the file directly to EFS in EC2 system

AWS Transfer Family now also supports file transfers to Amazon Elastic File System (Amazon EFS) file systems as well as Amazon S3.

upvoted 8 times

BBR01 Highly Voted 1 year, 1 month ago

Selected Answer: A

A should be enough. EFS can be mounted to ASG directly, and there is no need to use S3 in the middle.
upvoted 5 times

zdi561 Most Recent 3 months, 4 weeks ago

Selected Answer: A

A and D both work , but EFS is more natural file system for original batch routing, and could be faster than S3
upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: A

"... an on-premises file system receives" = file system > EFS.
SFT in AWS = AWS Transfer family

Even though D works, that would require changes in the current architecture.
upvoted 2 times

Scheldon 11 months, 2 weeks ago

Selected Answer: A

AnswerA

Option A and D will work, but taking into consideration requirements I would go with A
upvoted 2 times

Franjly 12 months ago

Selected Answer: A

File system = efs, fsx, ...
upvoted 2 times

7fb06b3 1 year ago

Selected Answer: D

Option D, I'm not 100% sure.. Always prefer S3 over EFS
upvoted 2 times

MatAlves 8 months, 3 weeks ago

"... an on-premises file system receives" = file system > EFS.
upvoted 2 times

JackyCCK 1 year, 2 months ago

I think the ans is A as well, option D require "Modify the application" which is not "minimize operational effort"
upvoted 4 times

khoantd 1 year, 2 months ago

Selected Answer: C

Option D
upvoted 1 times

Ipergorta 1 year, 2 months ago

Option D
upvoted 1 times

PolarFox 1 year, 3 months ago

Selected Answer: D

trasnfer + S3 = HA, scheduled scaling = resilient
upvoted 3 times

upvoted 6 times

NayeraB 1 year, 3 months ago

Selected Answer: D

I'm not 100% sure, but D looks like the right flow to me

upvoted 1 times

TwinSpark 1 year ago

I agree on a professional level, that's will make the company save money, that is the only things company care about. But for the exam i will go for A

upvoted 2 times

osmk 1 year, 3 months ago

Selected Answer: A

The service is designed to be highly scalable, highly available, and highly durable. Amazon EFS offers the following file system types to meet your availability and durability needs
->https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html
Amazon S3 achieves high availability by replicating data across multiple servers within AWS data centers->https://docs.aws.amazon.com/AmazonS3/latest/userguide/Welcome.html

upvoted 2 times

NayeraB 1 year, 3 months ago

But option A doesn't address the need for the application to pull the batch jobs from the new storage, also is the use of EFS needed here? In terms of it being a shared storage and whatnot..

upvoted 2 times

osmk 1 year, 3 months ago

A>>>>>

upvoted 1 times

osmk 1 year, 3 months ago

The service is designed to be highly scalable, highly available, and highly durable. Amazon EFS offers the following file system types to meet your availability and durability needs
->https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html
Amazon S3 achieves high availability by replicating data across multiple servers within AWS data centers->https://docs.aws.amazon.com/AmazonS3/latest/userguide/Welcome.html

upvoted 3 times

Andy_09 1 year, 4 months ago

Option D

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #754 *Topic 1*

A company has users all around the world accessing its HTTP-based application deployed on Amazon EC2 instances in multiple AWS Regions. The company wants to improve the availability and performance of the application. The company also wants to protect the application against common web exploits that may affect availability, compromise security, or consume excessive resources. Static IP addresses are required.

What should a solutions architect recommend to accomplish this?

A. Put the EC2 instances behind Network Load Balancers (NLBs) in each Region. Deploy AWS WAF on the NLBs. Create an accelerator using AWS Global Accelerator and register the NLBs as endpoints.

B. Put the EC2 instances behind Application Load Balancers (ALBs) in each Region. Deploy AWS WAF on the ALBs. Create an accelerator using AWS Global Accelerator and register the ALBs as endpoints. Most Voted

C. Put the EC2 instances behind Network Load Balancers (NLBs) in each Region. Deploy AWS WAF on the NLBs. Create an Amazon CloudFront distribution with an origin that uses Amazon Route 53 latency-based routing to route requests to the NLBs.

D. Put the EC2 instances behind Application Load Balancers (ALBs) in each Region. Create an Amazon CloudFront distribution with an origin that uses Amazon Route 53 latency-based routing to route requests to the ALBs. Deploy AWS WAF on the CloudFront distribution.

Correct Answer: *B*

Community vote distribution

B (86%) D (14%)

Comments

ogerber Highly Voted 9 months, 3 weeks ago

Selected Answer: B

HTTP based application so ALB is required.
because static IP addresses are required, we should use global accelerator:
"By default, Global Accelerator provides you with static IP addresses that you associate with your accelerator."

upvoted 12 times

Andy_09 Highly Voted 10 months, 1 week ago

Option D
upvoted 5 times

Typewriter101 9 months, 3 weeks ago

Why D cause i think global accelerator will do a better job an cloudfront to increase availability and performance
upvoted 2 times

Typewriter101 9 months, 3 weeks ago

than cloudfront*
upvoted 1 times

f07ed8f Most Recent 6 months, 2 weeks ago

Selected Answer: B

CloudFront doesn't support assigning a static IP address to distributions
upvoted 3 times

zinabu 7 months, 2 weeks ago

Selected Answer: B

Http based app=ALB
static IP= AWS global accelerator
for those who choice "A" NLB doesn't support Http based traffic it is just used for TCP/UDP based traffic.
upvoted 4 times

mohammadthainat 8 months, 1 week ago

Selected Answer: D

Something wrong in the question, here is why:

Static IP addresses are required --> NLB
protect against common web exploits --> WAF (But you can't use WAF directly with NLB)
HTTP-based application --> Cloudfront (using CloudFront with NLB is not recommended)
EC2s in multiple AWS Regions --> Route 53 latency-based
upvoted 3 times

mohammadthainat 8 months, 1 week ago

Changing my answer to B

Static IP addresses are required --> We can use Global Accelerator for fixed IP and WAF on the ALB
upvoted 4 times

TruthWS 8 months, 2 weeks ago

Option A
Static IP --> NLB
against common web exploits --> WAF
performance --> Global Accelerator is best choice in this situation.
upvoted 1 times

dkw2342 8 months, 2 weeks ago

No, option B is correct.

* WAF (L7) does not work with NLB (L4)
* Traffic enters via the Global Accelerator, so that's the customer-facing (static) IP - https://docs.aws.amazon.com/global-accelerator/latest/dg/about-accelerators.eip-accelerator.html
upvoted 3 times

jackky3123213 8 months, 2 weeks ago

Selected Answer: D

Option D
upvoted 1 times

alawada 8 months, 2 weeks ago

alawada 8 months, 2 weeks ago

Selected Answer: B

CloudFront uses multiple sets of dynamically changing IP addresses while Global Accelerator will provide you a set of static IP addresses as a fixed entry point to your applications

upvoted 2 times

lpergorta 8 months, 3 weeks ago

Option D

upvoted 1 times

Naveena_Devanga 9 months, 3 weeks ago

Correct Answer is C.
Static IP addresses are required specific to the requirement.

upvoted 1 times

stephensimudemy 9 months, 3 weeks ago

Selected Answer: B

CloudFront uses multiple sets of dynamically changing IP addresses while Global Accelerator will provide you a set of static IP addresses as a fixed entry point to your applications

upvoted 2 times

osmk 9 months, 3 weeks ago

Selected Answer: B

Network Load Balancer (NLB): NLB operates at layer 4 and does not support AWS WAF directly
https://docs.aws.amazon.com/elasticloadbalancing/latest/application/introduction.html

upvoted 2 times

osmk 9 months, 3 weeks ago

The company wants to improve the availability and performance of the application

upvoted 1 times

jaswantn 9 months, 3 weeks ago

Static IP addresses are required, so option B....global accelerator with ALB

upvoted 3 times

Dhokal 10 months ago

B is correct

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #755 *Topic 1*

A company's data platform uses an Amazon Aurora MySQL database. The database has multiple read replicas and multiple DB instances across different Availability Zones. Users have recently reported errors from the database that indicate that there are too many connections. The company wants to reduce the failover time by 20% when a read replica is promoted to primary writer.

Which solution will meet this requirement?

A. Switch from Aurora to Amazon RDS with Multi-AZ cluster deployment.

B. Use Amazon RDS Proxy in front of the Aurora database. **Most Voted**

C. Switch to Amazon DynamoDB with DynamoDB Accelerator (DAX) for read connections.

D. Switch to Amazon Redshift with relocation capability.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

osmk Highly Voted 9 months, 3 weeks ago

Selected Answer: B

By using Amazon RDS Proxy, your applications can pool and share database connections. This pooling improves scalability by allowing multiple application instances to reuse existing connections.
It also makes your applications more resilient to database failures. When a primary database instance fails, RDS Proxy automatically connects to a standby DB instance while preserving application connections.
=>https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html

upvoted 10 times

Umuntu Most Recent 10 months ago

Option B

upvoted 4 times

Andy_09 10 months, 1 week ago

Option B

upvoted 4 times

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #756 *Topic 1*

A company stores text files in Amazon S3. The text files include customer chat messages, date and time information, and customer personally identifiable information (PII).

The company needs a solution to provide samples of the conversations to an external service provider for quality control. The external service provider needs to randomly pick sample conversations up to the most recent conversation. The company must not share the customer PII with the external service provider. The solution must scale when the number of customer conversations increases.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Object Lambda Access Point. Create an AWS Lambda function that redacts the PII when the function reads the file. Instruct the external service provider to access the Object Lambda Access Point. Most Voted

B. Create a batch process on an Amazon EC2 instance that regularly reads all new files, redacts the PII from the files, and writes the redacted files to a different S3 bucket. Instruct the external service provider to access the bucket that does not contain the PII.

B. Create a web application on an Amazon EC2 instance that presents a list of the files, redacts the PII from the files, and allows the external service provider to download new versions of the files that have the PII redacted.

D. Create an Amazon DynamoDB table. Create an AWS Lambda function that reads only the data in the files that does not contain PII. Configure the Lambda function to store the non-PII data in the DynamoDB table when a new file is written to Amazon S3. Grant the external service provider access to the DynamoDB table.

Correct Answer: *A*

Community vote distribution

A (94%) D (6%)

Comments

osmk Highly Voted 1 year, 3 months ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonS3/latest/userguide/tutorial-s3-object-lambda-redact-pii.html

upvoted 13 times

Sergiuss95 1 year, 1 month ago

thanks
upvoted 2 times

Andy_09 Highly Voted 1 year, 4 months ago

Option A
upvoted 5 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: A

If you want to achieve the goal stipulated in the question, all 4 options work, but only option A minimizes operational overhead and ensures scalability while redacting PII dynamically.
upvoted 1 times

trinh_le 11 months, 2 weeks ago

Selected Answer: A

Use AWS Lambda functions to change the Object before it is retrieved by the caller application. Only one S3 bucket is needed, on top of which we create S3 Access Point And S3 Object Lambda Access Points
Use case:
1. Redact PII for analytics or non-production environment
2. Convert across data formats ex: XML to Json
3. Resizing and watermaking images on fly using caller-specific details ex: user who requested the object
upvoted 3 times

zinabu 1 year, 1 month ago

Selected Answer: D

Creating a RAID 0 array allows you to achieve a higher level of performance for a file system than you can provision on a single Amazon EBS volume. Use RAID 0 when I/O performance is of the utmost importance. With RAID 0, I/O is distributed across the volumes in a stripe.
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/raid-config.html
upvoted 1 times

Vlad 1 year, 3 months ago

A is the correct choice.
upvoted 3 times

Umuntu 1 year, 4 months ago

A is the best choice
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #757 *Topic 1*

A company is running a legacy system on an Amazon EC2 instance. The application code cannot be modified, and the system cannot run on more than one instance. A solutions architect must design a resilient solution that can improve the recovery time for the system.

What should the solutions architect recommend to meet these requirements?

A. Enable termination protection for the EC2 instance.

B. Configure the EC2 instance for Multi-AZ deployment.

C. Create an Amazon CloudWatch alarm to recover the EC2 instance in case of failure. **Most Voted**

D. Launch the EC2 instance with two Amazon Elastic Block Store (Amazon EBS) volumes that use RAID configurations for storage redundancy.

Correct Answer: *C*

Community vote distribution

C (80%) D (20%)

Comments

[Removed] Highly Voted 1 year, 1 month ago

Selected Answer: C

A. Enable termination protection for the EC2 instance.
No. Termination protection is about avoid accidentally delete the instance

B. Configure the EC2 instance for Multi-AZ deployment.
No. Question says "cannot run on more than one instance"

C. Create an Amazon CloudWatch alarm to recover the EC2 instance in case of failure.
Yes. CloudWatch can be used to recover the instance:
https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/UsingAlarmActions.html#AddingRecoverActions

D. Launch the EC2 instance with two Amazon Elastic Block Store (Amazon EBS) volumes that use RAID configurations for storage redundancy.
No. Raid could be helpful to increase resilience, but does not help with "improve the recovery time"

upvoted 17 times

Andy_09 Highly Voted 1 year, 4 months ago

Option C
upvoted 8 times

Typewriter101 1 year, 3 months ago

i think D is the answer.
Cause the question asks for a resilient solution and EBS with RAID config can balance between the performance and redundancy. EBS can also help with faster launch.
upvoted 2 times

MatAlves 8 months, 3 weeks ago

But how does that "improve the recovery time for the system"?
upvoted 2 times

mavik 1 year, 3 months ago

Your solution can't resolve the problem
upvoted 3 times

buzzinmumbai Most Recent 1 year, 2 months ago

Option should be B .They are not asking about storage anywhere. In muti-AZ you application runs on the primary and the secondary is kept in sync.
upvoted 1 times

mohammadthainat 1 year, 2 months ago
Selected Answer: C

Question about """improve the recovery time for the system""" RAID improves data resilience, but won't recover the instance if the system itself fails. it's 100% C
upvoted 5 times

dkw2342 1 year, 2 months ago

Pretty sure option D is NOT correct.

> RAID 5 and RAID 6 are not recommended for Amazon EBS (...).
> RAID 1 is also not recommended for use with Amazon EBS.

https://docs.aws.amazon.com/ebs/latest/userguide/raid-config.html#raid-config-options
upvoted 3 times

Awsbeginner87 1 year, 2 months ago

So what is the answer?
upvoted 2 times

Salilgen 5 months, 1 week ago

IMO answer is B.
You could set Max desired capacity to 1 in autoscaling group then you always have one instance running
upvoted 1 times

Salilgen 5 months, 1 week ago

Sorry. Answer is C.
In multi-AZ deployment you must have at least one instance per AZ
upvoted 1 times

sandordini 1 year, 1 month ago

C: You can create an Amazon CloudWatch alarm that monitors an Amazon EC2 instance and automatically recovers the instance if it becomes impaired due to an underlying hardware failure or a problem that requires AWS involvement to repair. Terminated instances cannot be recovered. A recovered instance is identical to the original instance, including the instance ID, private IP addresses, Elastic IP addresses, and all instance metadata.
upvoted 3 times

haci 1 year, 2 months ago

For those who choose C, the question asks that "must design a resilient solution".. C may improve recovery time but it has nothing to do with resiliency.

upvoted 1 times

JackyCCK 1 year, 2 months ago

"resilient solution that can improve the recovery time for the system" , resiliency here means only

upvoted 2 times

mavik 1 year, 3 months ago

Option C

upvoted 2 times

stephensimudemy 1 year, 3 months ago

Selected Answer: C

Can only run 1 instance.
improve recovery time.

upvoted 2 times

stephensimudemy 1 year, 3 months ago

Option B.
Question never ask anything about storage.

upvoted 1 times

osmk 1 year, 3 months ago

Selected Answer: D

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/raid-config.html

upvoted 6 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #758 *Topic 1*

A company wants to deploy its containerized application workloads to a VPC across three Availability Zones. The company needs a solution that is highly available across Availability Zones. The solution must require minimal changes to the application.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use Amazon Elastic Container Service (Amazon ECS). Configure Amazon ECS Service Auto Scaling to use target tracking scaling. Set the minimum capacity to 3. Set the task placement strategy type to spread with an Availability Zone attribute. **Most Voted**

B. Use Amazon Elastic Kubernetes Service (Amazon EKS) self-managed nodes. Configure Application Auto Scaling to use target tracking scaling. Set the minimum capacity to 3.

C. Use Amazon EC2 Reserved Instances. Launch three EC2 instances in a spread placement group. Configure an Auto Scaling group to use target tracking scaling. Set the minimum capacity to 3.

D. Use an AWS Lambda function. Configure the Lambda function to connect to a VPC. Configure Application Auto Scaling to use Lambda as a scalable target. Set the minimum capacity to 3.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

osmk Highly Voted 1 year, 3 months ago

Selected Answer: A

Amazon EKS self-managed nodes require you to manually install and configure the Kubernetes node components, such as kubelet, kube-proxy, and Docker, on your Amazon EC2 instances. You also need to manage the security group, IAM role, and subnet for your node group. Amazon ECS handles these tasks for you when you use the Amazon EC2 launch type .

upvoted 9 times

EdricHoang 11 months, 2 weeks ago

but it requires minimum change in application. I believe when changing to ECS, its a huge change

upvoted 1 times

ajwksdldfgdsg Most Recent 9 months ago

Selected Answer: A

Containerized.. = ECS

upvoted 1 times

1dd 1 year, 3 months ago

why not lambda?

upvoted 1 times

Sergiuss95 1 year, 1 month ago

Containerized... The solution must require minimal changes to the application.

upvoted 1 times

khoahoang 1 year, 2 months ago

lamda dont have containerized

upvoted 2 times

Andy_09 1 year, 4 months ago

Option A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #759 *Topic 1*

A media company stores movies in Amazon S3. Each movie is stored in a single video file that ranges from 1 GB to 10 GB in size.

The company must be able to provide the streaming content of a movie within 5 minutes of a user purchase. There is higher demand for movies that are less than 20 years old than for movies that are more than 20 years old. The company wants to minimize hosting service costs based on demand.

Which solution will meet these requirements?

A. Store all media content in Amazon S3. Use S3 Lifecycle policies to move media data into the Infrequent Access tier when the demand for a movie decreases.

B. Store newer movie video files in S3 Standard. Store older movie video files in S3 Standard-infrequent Access (S3 Standard-IA). When a user orders an older movie, retrieve the video file by using standard retrieval. Most Voted

C. Store newer movie video files in S3 Intelligent-Tiering. Store older movie video files in S3 Glacier Flexible Retrieval. When a user orders an older movie, retrieve the video file by using expedited retrieval.

D. Store newer movie video files in S3 Standard. Store older movie video files in S3 Glacier Flexible Retrieval. When a user orders an older movie, retrieve the video file by using bulk retrieval.

Correct Answer: *B*

Community vote distribution

B (55%) C (36%) A (9%)

Comments

Freddie26 Highly Voted 1 year, 3 months ago

Technically, expedited retrieval for files is not guaranteed within 1-5 minutes for files larger than 250 MB+. See https://docs.aws.amazon.com/AmazonS3/latest/userguide/restoring-objects-retrieval-options.html.
upvoted 17 times

osmk Highly Voted 1 year, 3 months ago

Selected Answer: B

S3 Standard-IA is for data that is accessed less frequently, but requires rapid access when needed. S3 Standard-IA offers the

high durability, high throughput, and low latency of S3 Standard, with a low per GB storage price and per GB retrieval chargehttps://aws.amazon.com/s3/storage-classes/
upvoted 12 times

tch Most Recent 3 months ago

Selected Answer: C

S3 Standard-infrequent Access (S3 Standard-IA) do not come with retrieval options??

https://docs.aws.amazon.com/AmazonS3/latest/userguide/restoring-objects-retrieval-options.html
upvoted 1 times

zdi561 3 months, 4 weeks ago

Selected Answer: C

Glacier flex expedited retrieval takes 1-5 min.
upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

Selected Answer: B

S3 Glacier Flexible Retrieval cannot grant the retrieval within 5 minutes even with expedited retrieval. Intellgent-Tiering is for unknow pattern, however, we already know higher demand for movies that are less than 20 years old than for movies that are more than 20 years old.
upvoted 1 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: B

A - Cost optimization for > 20-yr-old files is insufficient.
B - Good. Although, S3 Standard-IA is still relatively expensive compared to archival classes.
C - Glacier Flexible Retrieval may not consistently meet the 5-minute streaming requirement. Could be minutes, could be hours.
D - Bulk retrieval typically takes hours, let alone 5 minutes.
upvoted 3 times

tch 2 months, 2 weeks ago

this is the reason we choose C as answer! S3 Glacier Flexible Retrieval has option with 5 minutes time called expedited retrieval.
S3 Standard-IA do not come with option to retrieve data.... please check .....
https://aws.amazon.com/s3/storage-classes/glacier/
upvoted 1 times

Arad 6 months, 1 week ago

Selected Answer: C

C is the correct answer.
upvoted 1 times

Abdullah2004 9 months, 3 weeks ago

A is most correct □
upvoted 2 times

EdricHoang 11 months, 2 weeks ago

Selected Answer: B

"Expedited Retrieval, you can retrieve small amounts of data (up to 250 MB per request) within 1-5 minutes."
Cannot C
upvoted 8 times

NSA_Poker 11 months, 4 weeks ago

Selected Answer: B

(A) is eliminated. Demand metric is popularly based & cannot be configured with Lifecycle policies. Ex: an old movie can have resurgent demand 20 years after it's sequel is released.
(C) is eliminated. Expedited retrieval is for all but the largest archived objects (250 MB+).
(D) is eliminated. Bulk retrieval takes hours.
(B) is more expensive than S3 Glacier Flexible Retrieval but it's the only one that works.

(B) is more expensive than S3 Glacier Flexible Retrieval but it's the only one that works.
upvoted 6 times

bujuman 1 year, 1 month ago

Selected Answer: C

AS the pattern is uncertain- Customer could not, in advance, segregate data, the pattern will be determined on the fly - and with regard of the following S3 feature:
S3 Intelligent-Tiering is an additional storage class that provides flexibility for data with unknown or changing access patterns. It automates the movement of your objects between storage classes to optimize cost.
C will be the most cost effective for this use case.
upvoted 1 times

bujuman 1 year, 1 month ago

More insight:
S3 Glacier Flexible Retrieval for the most flexible retrieval options that balance cost with access times ranging from minutes to hours. Your retrieval options permit you to access all the archives you need, when you need them, for one low storage price.
This storage class comes with multiple retrieval options:
- Expedited retrievals (restore in 1–5 minutes)
- Standard retrievals (restore in 3–5 hours)
- Bulk retrievals (restore in 5–12 hours). Bulk retrievals are available at no additional charge
upvoted 3 times

Sergiuss95 1 year, 1 month ago

Selected Answer: C

Expedited 1-5min and for new files intelligent tier is a good option
upvoted 2 times

mohammadthainat 1 year, 2 months ago

Selected Answer: C

All old files should be in--> Glacier Flexible Retrieval takes (1-5 minutes) to retrieve the file.
New files should not stay in Standard Storage class forever --> Intelligent-Tiering
upvoted 3 times

JackyCCK 1 year, 2 months ago

I don't think C is an option, S3 Glacier Flexible takes hour to retrieve the data.
Option A is actually valid, but the way the option A describe it does not consider "demand patterns based on time"

So it should be B
upvoted 2 times

JackyCCK 1 year, 2 months ago

expedited retrieval should not be used in that way as well
upvoted 1 times

Drew3000 1 year, 2 months ago

Selected Answer: A

There is something I like about option A. It's the only one that deals with what happens with a movie that goes from "new" to "old". With other options, new movies will be new forever.
upvoted 2 times

dkw2342 1 year, 2 months ago

Option B makes the most sense.

Why not option C:

1. This is not an archival use case, the company runs a video streaming service, so objects are still accessed regularly. Accelerated Retrieval is designed for "occasional urgent requests for a subset of archives".
2. The 5 minute timeframe does not apply to items of 250+ MB.
3. Even if the timeframe were valid, it's not guaranteed ("typically")
4. Expedited retrieval is expensive if used frequently ($10.00 per 1,000 requests) - depending on access patterns, this may more than offset the savings in storage costs.

https://docs.aws.amazon.com/AmazonS3/latest/userguide/restoring-objects-retrieval-options.html
upvoted 4 times

TruthWS 1 year, 2 months ago

Option C
upvoted 1 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #760 *Topic 1*

A solutions architect needs to design the architecture for an application that a vendor provides as a Docker container image. The container needs 50 GB of storage available for temporary files. The infrastructure must be serverless.

Which solution meets these requirements with the LEAST operational overhead?

A. Create an AWS Lambda function that uses the Docker container image with an Amazon S3 mounted volume that has more than 50 GB of space.

B. Create an AWS Lambda function that uses the Docker container image with an Amazon Elastic Block Store (Amazon EBS) volume that has more than 50 GB of space.

C. Create an Amazon Elastic Container Service (Amazon ECS) cluster that uses the AWS Fargate launch type. Create a task definition for the container image with an Amazon Elastic File System (Amazon EFS) volume. Create a service with that task definition. **Most Voted**

D. Create an Amazon Elastic Container Service (Amazon ECS) cluster that uses the Amazon EC2 launch type with an Amazon Elastic Block Store (Amazon EBS) volume that has more than 50 GB of space. Create a task definition for the container image. Create a service with that task definition.

Correct Answer: *C*

Community vote distribution

C (94%) B (6%)

Comments

stephensimudemy Highly Voted 1 year, 3 months ago

Selected Answer: C

Options A and B involve AWS Lambda, which is suitable for event-driven, short-lived compute tasks, but it's NOT ideal for long-running containerized applications and managing large volumes of data.

upvoted 8 times

Andy_09 Highly Voted 1 year, 4 months ago

Option C

upvoted 6 times

nj1999 1 year, 4 months ago

Why C and not B?
upvoted 1 times

03beafc 1 year, 1 month ago

Not B because your lambda container needs the RIC and the image is already provided, presumably without the RIC (or else it would have mentioned it)
upvoted 3 times

sandordini 1 year, 1 month ago

RIC: Runtime interface clients
upvoted 2 times

hajra313 1 year, 3 months ago

the infrastructure must be serverless
upvoted 2 times

Cali182 1 year, 3 months ago

Creating an AWS Lambda function that uses the Docker container image with an Amazon S3 mounted volume might not be suitable because Lambda functions have limitations on execution duration (15 minutes) and storage size (maximum 512 MB in the /tmp directory).
upvoted 5 times

dkw2342 1 year, 2 months ago

There's no indication of runtime, so that's not the reason.

A is wrong because "S3 volumes" do not exist. If the question were about S3 buckets: while it is possible to mount an S3 bucket using FUSE, this is completely unsupported by AWS and definitely won't work in a container running on Lambda (you can't assign SYS_ADMIN cap and mount /dev/fuse).

B is wrong because you can't use EBS volumes with Lambda.

As an aside, Lambda supports up to 10 GB of ephemeral storage (configurable).
upvoted 3 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: C

A - Technically, temporary file access in Lambda is limited to the /tmp directory with a 512 MB limit.
B - AWS Lambda does not natively support mounting EBS volumes.
C - Looks good to me.
D - ECS with EC2 is not serverless.

FYI - Serverless infrastructure: Solutions must not involve managing servers or underlying infrastructure. AWS Fargate and AWS Lambda are serverless options, while Amazon ECS with EC2 launch type involves managing EC2 instances.
upvoted 1 times

dragongoseki 11 months, 3 weeks ago

Selected Answer: C

C is right answer.
upvoted 2 times

DZRomero 11 months, 3 weeks ago

Selected Answer: C

The combination of ECS with Fargate and EFS (option C) provides a serverless solution that can run Docker containers and meet the storage requirements, all while minimizing operational overhead. You don't need to manage any servers, and the storage will automatically scale as needed. This makes it the best fit for the given requirements.
upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: C

Lambda would need Runtime interface clients (RIC) to host a container workload.
Also Lambda storage limit: 10GB
Fargate is Serverless >> C

upvoted 4 times

zinabu 1 year, 1 month ago

Selected Answer: B

the key word here is {" serverless + temporary file"}
A: it uses S3 for storage that is not a temporary file storage system
C: that was good using ECS with farget for serverless part but is uses EFS file system still it is a durable file system not temporary
D: Using EBS was good to use for temporary file system but it is mounted on EC2 which is not serverless. so that we are left with "B" which uses [lambda(serverless) + EBS(temporary storage)]

upvoted 1 times

zinabu 1 year, 1 month ago

the key word here is {" serverless + temporary file"}
A: it uses S3 for storage that is not a temporary file storage system
C: that was good using ECS with farget for serverless part but is uses EFS file system still it is a durable file system not temporary
D: Using EBS was good to use for temporary file system but it is mounted on EC2 which is not serverless. so that we are left with "B" which uses [lambda(serverless) + EBS(temporary storage)]

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #761 *Topic 1*

A company needs to use its on-premises LDAP directory service to authenticate its users to the AWS Management Console. The directory service is not compatible with Security Assertion Markup Language (SAML).

Which solution meets these requirements?

A. Enable AWS IAM Identity Center (AWS Single Sign-On) between AWS and the on-premises LDAP.

B. Create an IAM policy that uses AWS credentials, and integrate the policy into LDAP.

C. Set up a process that rotates the IAM credentials whenever LDAP credentials are updated.

D. Develop an on-premises custom identity broker application or process that uses AWS Security Token Service (AWS STS) to get short-lived credentials. Most Voted

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

kempes Highly Voted 1 year, 4 months ago

Selected Answer: D

The solution that best meets the requirements. This approach provides a pathway for authenticating LDAP users to AWS without requiring direct LDAP to AWS IAM Identity Center integration or SAML compatibility, offering a flexible and secure method to extend on-premises authentication mechanisms to AWS services.

upvoted 11 times

aditianand 1 year ago

Why not option A A. Enable AWS IAM Identity Center (AWS Single Sign-On) between AWS and the on-premises LDAP.

upvoted 1 times

NSA_Poker 11 months, 4 weeks ago

(A) is incorrect bc to use AWS IAM Identity Center (AWS Single Sign-On) with an external IdP, you need SAML.

upvoted 5 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: D

A - AWS IAM Identity Center (formerly AWS SSO) allows integration with SAML-compatible identity providers, which won't work with on-premises LDAP directory (not SAML-compatible).
B - There's no way to "integrate" an IAM policy into LDAP.
C - Too complex. AWS already provides better mechanisms for handling temporary credentials - AWS STS.
D - YES. Actually, this approach is commonly used when SAML is not an option.

upvoted 1 times

Scheldon 11 months, 3 weeks ago

Selected Answer: D

AnswerD

upvoted 2 times

ike001 11 months, 3 weeks ago

D is the answer

upvoted 2 times

NSA_Poker 11 months, 4 weeks ago

Selected Answer: D

Identity federation can be accomplished in one of three ways.
(1) Use a corporate IdP (such as Microsoft Active Directory) or a custom identity broker application. Each option uses AWS STS.
(2) Create an integration that uses Security Assertion Markup Language (SAML).
(3) Use a web identity provider, such as Amazon Cognito.

upvoted 2 times

1e22522 10 months ago

YEA SURE FED

upvoted 1 times

TwinSpark 1 year ago

Selected Answer: D

option D
As per described here:
https://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_common-scenarios_federated-users.html#id_roles_common-scenarios_federated-users-idbroker
option A is wrong becouse for use SSO need to be compatible with SAML (at least this is what i understand from here: https://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_common-scenarios_federated-users.html#id_roles_common-scenarios_federated-users-saml20)

upvoted 2 times

Naveena_Devanga 1 year, 3 months ago

Option D
A custom identity broker application can be built to perform a similar function to an identity store that is not compatible with SAML 2.0. The broker application authenticates users, requests temporary credentials from AWS, and provides them to the user to access AWS resources.

upvoted 2 times

aditianand 1 year ago

Why not option A A. Enable AWS IAM Identity Center (AWS Single Sign-On) between AWS and the on-premises LDAP.

upvoted 1 times

jaswantn 1 year, 3 months ago

If your identity store is not compatible with SAML 2.0, then you can build a custom identity broker application to perform a similar function. .....option D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #762 *Topic 1*

A company stores multiple Amazon Machine Images (AMIs) in an AWS account to launch its Amazon EC2 instances. The AMIs contain critical data and configurations that are necessary for the company's operations. The company wants to implement a solution that will recover accidentally deleted AMIs quickly and efficiently.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create Amazon Elastic Block Store (Amazon EBS) snapshots of the AMIs. Store the snapshots in a separate AWS account.

B. Copy all AMIs to another AWS account periodically.

C. Create a retention rule in Recycle Bin. **Most Voted**

D. Upload the AMIs to an Amazon S3 bucket that has Cross-Region Replication.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

alawada Highly Voted 8 months, 2 weeks ago

Selected Answer: C

Recycle Bin is a data recovery feature that enables you to restore accidentally deleted Amazon EBS snapshots and EBS-backed AMIs. When using Recycle Bin, if your resources are deleted, they are retained in the Recycle Bin for a time period that you specify before being permanently deleted. You can restore a resource from the Recycle Bin at any time before its retention period expires. This solution has the least operational overhead, as you do not need to create, copy, or upload any additional resources. You can also manage tags and permissions for AMIs in the Recycle Bin. AMIs in the Recycle Bin do not incur any additional charges.Reference:

upvoted 8 times

[Removed] Most Recent 8 months, 2 weeks ago

Selected Answer: C

Option C
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/recycle-bin-working-with-rules.html

upvoted 2 times

asdfcdsxdfc 9 months, 1 week ago

[illegible] 9 months, 1 week ago

Selected Answer: C

C is correct

upvoted 2 times

Naveena_Devanga 9 months, 3 weeks ago

Option C
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/recycle-bin-working-with-rules.html

upvoted 2 times

Freddie26 9 months, 3 weeks ago

Option C is correct. Recycling bin is a new feature to protect snaps and AMIs from accidental or malicious deleting. Inside the recycling bin, set a retention policy, and then your images or snapshots are protected.

upvoted 4 times

mestule 10 months ago

Selected Answer: C

https://aws.amazon.com/about-aws/whats-new/2022/02/amazon-ec2-recycle-bin-machine-images/

upvoted 4 times

Andy_09 10 months, 1 week ago

Option C

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #763 *Topic 1*

A company has 150 TB of archived image data stored on-premises that needs to be moved to the AWS Cloud within the next month. The company's current network connection allows up to 100 Mbps uploads for this purpose during the night only.

What is the MOST cost-effective mechanism to move this data and meet the migration deadline?

A. Use AWS Snowmobile to ship the data to AWS.

B. Order multiple AWS Snowball devices to ship the data to AWS. **Most Voted**

C. Enable Amazon S3 Transfer Acceleration and securely upload the data.

D. Create an Amazon S3 VPC endpoint and establish a VPN to upload the data.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option B

upvoted 13 times

Scheldon Most Recent 11 months, 3 weeks ago

Selected Answer: B

AnwerB

SnowBall Storage will give us 80TB. To transfer data we will need 2 devices. Taking into consideration that 1 device = ~300$ we will spend 600$. Option B is most Cost effective and will allow us to end operation in lest then month.

upvoted 3 times

sandordini 1 year, 1 month ago

Selected Answer: B

Snowball base free from 200USD, Snowmobile base fee from 4100USD (According to AWS)

upvoted 4 times

sandordini 1 year, 1 month ago

Snowmobile advised above 10 Petabytes
Snowball(s) below 10 PB
upvoted 3 times

TruthWS 1 year, 2 months ago

Option B - Snowmobile have higher cost
upvoted 2 times

Mikado211 1 year, 2 months ago

Selected Answer: B

Amazon S3 Transfer Acceleration must be very expensive
Correct in such case : B Snowball
upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: B

B is correct
upvoted 2 times

iczcezar 1 year, 3 months ago

Option B
upvoted 2 times

Naveena_Devanga 1 year, 3 months ago

Option B:
1 Snow Ball Max Allowed capacity is 80 TB. Hence, you need to order multiple snowballs to achieve the requirement.
upvoted 3 times

stephensimudemy 1 year, 3 months ago

Selected Answer: B

B. Its only 150TB
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #764 *Topic 1*

A company wants to migrate its three-tier application from on premises to AWS. The web tier and the application tier are running on third-party virtual machines (VMs). The database tier is running on MySQL.

The company needs to migrate the application by making the fewest possible changes to the architecture. The company also needs a database solution that can restore data to a specific point in time.

Which solution will meet these requirements with the LEAST operational overhead?

A. Migrate the web tier and the application tier to Amazon EC2 instances in private subnets. Migrate the database tier to Amazon RDS for MySQL in private subnets.

B. Migrate the web tier to Amazon EC2 instances in public subnets. Migrate the application tier to EC2 instances in private subnets. Migrate the database tier to Amazon Aurora MySQL in private subnets. **Most Voted**

C. Migrate the web tier to Amazon EC2 instances in public subnets. Migrate the application tier to EC2 instances in private subnets. Migrate the database tier to Amazon RDS for MySQL in private subnets.

D. Migrate the web tier and the application tier to Amazon EC2 instances in public subnets. Migrate the database tier to Amazon Aurora MySQL in public subnets.

Correct Answer: *B*

Community vote distribution

B (92%) C (8%)

Comments

haci Highly Voted 1 year, 3 months ago

Selected Answer: B

I'm between B and C. Since RDS requires an additional configuration for PTR, it adds an operational overhead. So I will go with B.

Aurora provides automated backup and point-in-time recovery, simplifying backup management and data protection. Continuous incremental backups are taken automatically and stored in Amazon S3, and data retention periods can be specified to meet compliance requirements.

RDS provides the same but first, the users should set a retention period for these backups, allowing historical data recovery in case of accidental data loss or corruption, and point-in-time recovery (PITR) allows users to restore the database to any specific moment within the set retention period.

upvoted 13 times

Bwhizzy Most Recent 8 months, 1 week ago

Selected Answer: B

Answer is B. please see below article
https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-pitr.html

upvoted 2 times

kbgsgsgs 8 months, 2 weeks ago

Selected Answer: C

Auroralo maigeuleisyeonhalyeomyeon MySQLboda akitegcheo byeongyeong-i deo pil-yohal su issseubnida. ganeunghamyeon choesohan-ui byeongyeong-eul yocheong
67 / 5,000
Migrating to Aurora may require more architectural changes than MySQL. Request minimal changes if possible.

upvoted 1 times

Scheldon 11 months, 3 weeks ago

Selected Answer: B

AnswerB

The problem with this question is that we do not have enough information. We can execute task with both AURORA or RDS DB. I will go with AURORA as it is Amazon Proprietary and is developed by AWS teams, hence we do not need to think about updates etc. as it is done by AWS teams.

upvoted 2 times

MattBJ 1 year, 2 months ago

Selected Answer: B

B is the correct option.

upvoted 2 times

shahreh1 1 year, 3 months ago

B: Amazon Aurora is a fully managed relational database engine that's compatible with both MySQL and PostgreSQL

upvoted 4 times

DEN_ZZ 1 year, 3 months ago

Selected Answer: B

PTR, it's Aurora

upvoted 4 times

stephensimudemy 1 year, 3 months ago

Selected Answer: C

It's C. Strictly speaking, there is no AWS DB call Amazon Aurora "MySQL"

upvoted 1 times

ogerber 1 year, 3 months ago

https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.AuroraMySQL.html

upvoted 3 times

hajra313 1 year, 4 months ago

C. This option aligns with the requirements by keeping the web tier in public subnets, migrating the application tier to EC2 instances in private subnets to enhance security, and using Amazon RDS for MySQL in private subnets to meet the database requirements with minimal operational overhead. option A:While migrating the web tier and application tier to EC2 instances in private subnets minimizes exposure to the internet. option B:. Migrating the database tier to Amazon Aurora MySQL introduces changes to the database engine, which might require additional testing and adjustments to the application. Additionally, Aurora MySQL does not directly support point-in-time recovery; instead, it uses continuous backups and snapshots for data recovery.

upvoted 3 times

Andy_09 1 year, 4 months ago

Andy_09 1 year, 4 months ago

Option A works better

upvoted 2 times

Andy_09 1 year, 4 months ago

Option B

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #765 *Topic 1*

A development team is collaborating with another company to create an integrated product. The other company needs to access an Amazon Simple Queue Service (Amazon SQS) queue that is contained in the development team's account. The other company wants to poll the queue without giving up its own account permissions to do so.

How should a solutions architect provide access to the SQS queue?

A. Create an instance profile that provides the other company access to the SQS queue.

B. Create an IAM policy that provides the other company access to the SQS queue.

C. Create an SQS access policy that provides the other company access to the SQS queue. **Most Voted**

D. Create an Amazon Simple Notification Service (Amazon SNS) access policy that provides the other company access to the SQS queue.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

iczcezar Highly Voted 1 year, 3 months ago

The correct option to provide access to the SQS queue without giving up the other company's account permissions is:

C. Create an SQS access policy that provides the other company access to the SQS queue.

By creating an SQS access policy, you can define specific permissions for the other company to access the SQS queue without requiring them to modify their own account permissions. This allows for fine-grained control over access to the queue while maintaining security and isolation between accounts. Options A, B, and D are not appropriate for granting access to the SQS queue in this scenario.

upvoted 8 times

7fb06b3 Highly Voted 1 year ago

Selected Answer: C

Amazon SQS policy system lets you grant permission to other Amazon Accounts.
https://docs.amazonaws.cn/en_us/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-using-identity-based-policies.html

upvoted 5 times

FlyingHawk Most Recent 5 months, 1 week ago

Selected Answer: C

https://repost.aws/knowledge-center/sqs-queue-access-permissions
Here is the example:
{
"Version": "2012-10-17",
"Statement": [
{
"Effect": "Allow",
"Principal": {
"AWS": "arn:aws:iam::OtherCompanyAccountID:root"
},
"Action": [
"sqs:ReceiveMessage",
"sqs:DeleteMessage"
],
"Resource": "arn:aws:sqs:region:YourAccountID:YourQueueName"
}
]
}

upvoted 1 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: C

A - Instance profiles are attached to Amazon EC2 instances. This is irrelevant.
B - While you could create an IAM role in your account for the other company to assume, it is explicitly stated that the other company does not want to give up its own permissions, meaning cross-account access must be configured, but IAM policies alone do not provide cross-account access.
C - SQS supports resource-based policies, which allow you to grant access to specific AWS principals (users, roles, or accounts) across accounts.
D - No idea what it would achieve. SNS policies don't even apply to SQS queues.

upvoted 2 times

Scheldon 11 months, 3 weeks ago

Selected Answer: C

AnswerC

Creating Access policy in SQS which will allow other company to acess SQS queue seems to be the only solution which is RIGHT here

upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: C

SQS Access Policy for secure, fine-grained Cross-account access

upvoted 3 times

NayeraB 1 year, 3 months ago

Selected Answer: C

https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-overview-of-managing-access.html

upvoted 2 times

hajra313 1 year, 4 months ago

option A: Instance profiles are used to grant permissions to EC2 instances, not for granting access to other AWS services like SQS queues. Option B.:AM policies are applied to IAM users, groups, or roles within the same AWS account. They are not directly applicable to granting access to resources in other AWS accounts. option C:SQS access policies allow you to grant cross-account access to SQS resources. You can specify the necessary permissions in the policy and attach it directly to the SQS queue. This way, you can give the other company's AWS account the necessary permissions to poll the queue without compromising their account permissions. option D. Amazon SNS access policies are used to manage access to SNS topics, not SQS queues

upvoted 4 times

kempes 1 year, 4 months ago

Selected Answer: C

Option C
upvoted 4 times

Andy_09 1 year, 4 months ago

Option B
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #766 *Topic 1*

A company's developers want a secure way to gain SSH access on the company's Amazon EC2 instances that run the latest version of Amazon Linux. The developers work remotely and in the corporate office.

The company wants to use AWS services as a part of the solution. The EC2 instances are hosted in a VPC private subnet and access the internet through a NAT gateway that is deployed in a public subnet.

What should a solutions architect do to meet these requirements MOST cost-effectively?

A. Create a bastion host in the same subnet as the EC2 instances. Grant the ec2:CreateVpnConnection IAM permission to the developers. Install EC2 Instance Connect so that the developers can connect to the EC2 instances.

B. Create an AWS Site-to-Site VPN connection between the corporate network and the VPC. Instruct the developers to use the Site-to-Site VPN connection to access the EC2 instances when the developers are on the corporate network. Instruct the developers to set up another VPN connection for access when they work remotely.

C. Create a bastion host in the public subnet of the VPConfigure the security groups and SSH keys of the bastion host to only allow connections and SSH authentication from the developers' corporate and remote networks. Instruct the developers to connect through the bastion host by using SSH to reach the EC2 instances.

D. Attach the AmazonSSMManagedInstanceCore IAM policy to an IAM role that is associated with the EC2 instances. Instruct the developers to use AWS Systems Manager Session Manager to access the EC2 instances. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

alawada Highly Voted 8 months, 2 weeks ago

Selected Answer: D

AWS Systems Manager Session Manager is a service that enables you to securely connect to your EC2 instances without using SSH keys or bastion hosts. You can use Session Manager to access your instances through the AWS Management Console, the AWS CLI, or the AWS SDKs. Session Manager uses IAM policies and roles to control who can access which instances. By attaching the AmazonSSMManagedInstanceCore IAM policy to an IAM role that is associated with the EC2 instances, you grant the Session Manager service the necessary permissions to perform actions on your instances. You also need to attach another IAM policy to the developers' IAM users or roles that allows them to start sessions to the instances.

upvoted 7 times

Andy_09 Highly Voted 10 months, 1 week ago

Option D

upvoted 5 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: D

A - Bastion hosts are typically placed in public subnets (not private subnets). Also, see option C below.
B - Would work, but this is too complex and too expensive.
C - An additional EC2 instance (bastion host) = instance running costs + endless maintenance.
D - God yes. No need to manage SSH keys, and access is encrypted by default.

upvoted 1 times

TruthWS 8 months, 2 weeks ago

Option D

upvoted 3 times

Mikado211 8 months, 2 weeks ago

Selected Answer: D

SSM is always the recommended way of connection for EC2 "using ssh".
It's the most cost effective and the most secure way of doing the job.

upvoted 2 times

iczcezar 9 months, 2 weeks ago

Why not C?

upvoted 2 times

pila21 8 months, 3 weeks ago

it doesn`t meet requirements MOST cost-effectively

upvoted 3 times

kempes 10 months ago

Selected Answer: D

Option D

upvoted 3 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #767 *Topic 1*

A pharmaceutical company is developing a new drug. The volume of data that the company generates has grown exponentially over the past few months. The company's researchers regularly require a subset of the entire dataset to be immediately available with minimal lag. However, the entire dataset does not need to be accessed on a daily basis. All the data currently resides in on-premises storage arrays, and the company wants to reduce ongoing capital expenses.

Which storage solution should a solutions architect recommend to meet these requirements?

A. Run AWS DataSync as a scheduled cron job to migrate the data to an Amazon S3 bucket on an ongoing basis.

B. Deploy an AWS Storage Gateway file gateway with an Amazon S3 bucket as the target storage. Migrate the data to the Storage Gateway appliance.

C. Deploy an AWS Storage Gateway volume gateway with cached volumes with an Amazon S3 bucket as the target storage. Migrate the data to the Storage Gateway appliance. **Most Voted**

D. Configure an AWS Site-to-Site VPN connection from the on-premises environment to AWS. Migrate data to an Amazon Elastic File System (Amazon EFS) file system.

Correct Answer: *C*

Community vote distribution

C (91%) B (9%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option C

upvoted 11 times

hajra313 Highly Voted 1 year, 4 months ago

B. Deploying an AWS Storage Gateway file gateway with an Amazon S3 bucket as the target storage would require the entire dataset to be stored in Amazon S3, which might not be cost-effective considering that only a subset of the data needs to be accessed regularly. Additionally, accessing data directly from S3 might introduce latency. so correct option is C bcz AWS Storage Gateway volume gateway with cached volumes allows the company to keep frequently accessed data locally on-premises while storing the entire dataset in Amazon S3. This solution provides immediate access to the subset of data with minimal lag, as frequently accessed data is cached locally. It also reduces ongoing capital expenses as it leverages Amazon S3 storage, which is cost-effective.

upvoted 8 times

·

capepenguin 11 months, 1 week ago

Both B and C store the entire data set in Amazon S3 and the AWS Storage Gateway file gateway also supports local caching. I do not know if it is B or C for me.

upvoted 2 times

nadeerm Most Recent 3 months ago

Selected Answer: B

Volume Gateway is designed for block storage (EBS volumes), not file storage. It is not the best fit for this use case, which involves file-based data access.

upvoted 2 times

Scheldon 11 months, 2 weeks ago

Selected Answer: C

AnswerC

Deploying an AWS Storage Gateway volume gateway with cached volumes will allow to store all data in AWS but the most frequently accessed data will be stored/cached localy (on-premises) = low latency for most used data while all data will be stored in the cloud.

https://docs.aws.amazon.com/storagegateway/latest/vgw/StorageGatewayConcepts.html#storage-gateway-cached-concepts

upvoted 3 times

BatVanyo 1 year, 1 month ago

A storage guy here.. the question is not clear enough to give a definitive answer between B and C, as both can do the job.
An "on-prem storage array" can be any of the three:
- File array (serving any file protocol, e.g. NFS/SMB) -> requiring a file gateway (supports caching of the most recently used data)
- Block array (iSCSI/Fibre Channel) -> requiring a volume gateway (supports cached volumes most recently used data)
- Combo (providing both File and Block protocols)

Something is clearly missing in the question in order to give a definitive answer between B and C.

upvoted 4 times

mohammadthainat 1 year, 2 months ago

Selected Answer: C

storage arrays = Volume Gateway

upvoted 3 times

lenotc 1 year, 2 months ago

Selected Answer: C

storage array, also known as a disk array so AWS Storage Gateway volume.
its a trap

upvoted 5 times

MattBJ 1 year, 2 months ago

Selected Answer: C

C is correct. Using AWS Storage Gateway volume gateway with cached volumes provide local access to the file.

upvoted 3 times

ninasgx 1 year, 3 months ago

Selected Answer: C

require a subset of the entire dataset => cached volumes

upvoted 4 times

osmk 1 year, 3 months ago

Selected Answer: C

The company's researchers regularly require a subset of the entire dataset to be immediately available with minimal laghttps://docs.aws.amazon.com/storagegateway/latest/vgw/WhatIsStorageGateway.html

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #768 *Topic 1*

A company has a business-critical application that runs on Amazon EC2 instances. The application stores data in an Amazon DynamoDB table. The company must be able to revert the table to any point within the last 24 hours.

Which solution meets these requirements with the LEAST operational overhead?

A. Configure point-in-time recovery for the table. **Most Voted**

B. Use AWS Backup for the table.

C. Use an AWS Lambda function to make an on-demand backup of the table every hour.

D. Turn on streams on the table to capture a log of all changes to the table in the last 24 hours. Store a copy of the stream in an Amazon S3 bucket.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option A

upvoted 13 times

MattBJ Highly Voted 1 year, 2 months ago

Selected Answer: A

A is correct. One of the highlight features of DynamoDB.

upvoted 5 times

Scheldon Most Recent 11 months, 2 weeks ago

Selected Answer: A

AnswerA
Point-in-time recovery helps protect your DynamoDB tables from accidental write or delete operations. With point-in-time recovery, you don't have to worry about creating, maintaining, or scheduling on-demand backups. For example, suppose that a test script writes accidentally to a production DynamoDB table. With point-in-time recovery, you can restore that table to any point in time during the last 35 days. After you enable point-in-time recovery, you can restore to any point in time from five minutes before the current time until 35 days ago. DynamoDB maintains incremental backups of your table.

https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/PointInTimeRecovery.html
upvoted 3 times

1dd 1 year, 3 months ago

Selected Answer: A

option A
upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: A

A looks correct
upvoted 3 times

mavik 1 year, 3 months ago

Selected Answer: A

Option A
upvoted 3 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #769 *Topic 1*

A company hosts an application used to upload files to an Amazon S3 bucket. Once uploaded, the files are processed to extract metadata, which takes less than 5 seconds. The volume and frequency of the uploads varies from a few files each hour to hundreds of concurrent uploads. The company has asked a solutions architect to design a cost-effective architecture that will meet these requirements.

What should the solutions architect recommend?

A. Configure AWS CloudTrail trails to log S3 API calls. Use AWS AppSync to process the files.

B. Configure an object-created event notification within the S3 bucket to invoke an AWS Lambda function to process the files. **Most Voted**

C. Configure Amazon Kinesis Data Streams to process and send data to Amazon S3. Invoke an AWS Lambda function to process the files.

D. Configure an Amazon Simple Notification Service (Amazon SNS) topic to process the files uploaded to Amazon S3. Invoke an AWS Lambda function to process the files.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

hajra313 Highly Voted 1 year, 4 months ago

option b bcz option c is WS AppSync is not the most appropriate solution for file processing.
option d While Amazon Simple Notification Service (SNS) can be used to trigger actions based on S3 events, it's not directly involved in processing files .option c :Kinesis is typically used for real-time data streaming and analytics, which may not be needed for simple file processing tasks such as extracting metadata.

upvoted 5 times

mestule Highly Voted 1 year, 4 months ago

Selected Answer: B

B seems to be make most sense to me.

upvoted 5 times

Lin878 Most Recent 11 months, 2 weeks ago

Selected Answer: B

https://docs.aws.amazon.com/AmazonS3/latest/userguide/EventNotifications.html
upvoted 3 times

Scheldon 11 months, 2 weeks ago

Selected Answer: B

AnswerB
upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: B

The problem with C is how it sends the data to S3, if it was Firehose it would make sense. I waka for B.
upvoted 3 times

MattBJ 1 year, 2 months ago

Selected Answer: B

B is correct. The most cost effective option.
upvoted 3 times

jaswantn 1 year, 3 months ago

option B
upvoted 2 times

kempes 1 year, 4 months ago

Option D
upvoted 2 times

Andy_09 1 year, 4 months ago

Option D
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #770 *Topic 1*

A company's application is deployed on Amazon EC2 instances and uses AWS Lambda functions for an event-driven architecture. The company uses nonproduction development environments in a different AWS account to test new features before the company deploys the features to production.

The production instances show constant usage because of customers in different time zones. The company uses nonproduction instances only during business hours on weekdays. The company does not use the nonproduction instances on the weekends. The company wants to optimize the costs to run its application on AWS.

Which solution will meet these requirements MOST cost-effectively?

A. Use On-Demand Instances for the production instances. Use Dedicated Hosts for the nonproduction instances on weekends only.

B. Use Reserved Instances for the production instances and the nonproduction instances. Shut down the nonproduction instances when not in use.

C. Use Compute Savings Plans for the production instances. Use On-Demand Instances for the nonproduction instances. Shut down the nonproduction instances when not in use. **Most Voted**

D. Use Dedicated Hosts for the production instances. Use EC2 Instance Savings Plans for the nonproduction instances.

Correct Answer: *C*

Community vote distribution

C (87%) B (13%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option C

upvoted 13 times

MattBJ Highly Voted 1 year, 2 months ago

Selected Answer: C

Definitely C.

upvoted 5 times

Penjerla Most Recent 5 months, 2 weeks ago

Selected Answer: B

A Reserved Instance is preferred if you use the reserve capacity for the full duration of the one- or three-year term. You may also find that a Reserved Instance offers cost savings over On-Demand Instances, even if you are only using the instance at half capacity. Reserved Instances also offer guaranteed availability for stop-and-start scenarios. In addition, by utilizing and continuously monitoring Convertible Reserved Instances, you can save more money.

upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: C

"An Amazon EC2 Dedicated Host is a physical server fully dedicated for your use, so you can help address corporate compliance requirements."

https://aws.amazon.com/ec2/dedicated-hosts/

That already eliminates "A" and "D". No need to pay extra since there is no requirement for eligible software licences.

B is INCORRECT = no need to use reserved instances for non-prod, since they will only be active during business hours during weekdays.

upvoted 2 times

Scheldon 11 months, 2 weeks ago

Selected Answer: C

AnswerC

We need to use somekind of savings plans for production and disable Development servers when we are not using them. Only Option C has both.

upvoted 2 times

Naveena_Devanga 1 year, 3 months ago

Option C

upvoted 2 times

stephensimudemy 1 year, 3 months ago

Selected Answer: C

It's C

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #771 *Topic 1*

A company stores data in an on-premises Oracle relational database. The company needs to make the data available in Amazon Aurora PostgreSQL for analysis. The company uses an AWS Site-to-Site VPN connection to connect its on-premises network to AWS.

The company must capture the changes that occur to the source database during the migration to Aurora PostgreSQL.

Which solution will meet these requirements?

A. Use the AWS Schema Conversion Tool (AWS SCT) to convert the Oracle schema to Aurora PostgreSQL schema. Use the AWS Database Migration Service (AWS DMS) full-load migration task to migrate the data.

B. Use AWS DataSync to migrate the data to an Amazon S3 bucket. Import the S3 data to Aurora PostgreSQL by using the Aurora PostgreSQL aws_s3 extension.

C. Use the AWS Schema Conversion Tool (AWS SCT) to convert the Oracle schema to Aurora PostgreSQL schema. Use AWS Database Migration Service (AWS DMS) to migrate the existing data and replicate the ongoing changes. **Most Voted**

D. Use an AWS Snowball device to migrate the data to an Amazon S3 bucket. Import the S3 data to Aurora PostgreSQL by using the Aurora PostgreSQL aws_s3 extension.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

kempes Highly Voted 1 year, 4 months ago

Selected Answer: C

Option C

upvoted 10 times

Andy_09 Highly Voted 1 year, 4 months ago

Option C

upvoted 8 times

striker89 Most Recent 7 months, 1 week ago

Selected Answer: C

migrate the existing data and replicate the ongoing changes.
upvoted 2 times

Scheldon 11 months, 3 weeks ago

Selected Answer: C

Answer C
upvoted 1 times

MattBJ 1 year, 2 months ago

Selected Answer: C

C is correct. As we need to capture the change during the migration.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #772 *Topic 1*

A company built an application with Docker containers and needs to run the application in the AWS Cloud. The company wants to use a managed service to host the application.

The solution must scale in and out appropriately according to demand on the individual container services. The solution also must not result in additional operational overhead or infrastructure to manage.

Which solutions will meet these requirements? (Choose two.)

A. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate. **Most Voted**

B. Use Amazon Elastic Kubernetes Service (Amazon EKS) with AWS Fargate. **Most Voted**

C. Provision an Amazon API Gateway API. Connect the API to AWS Lambda to run the containers.

D. Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 worker nodes.

E. Use Amazon Elastic Kubernetes Service (Amazon EKS) with Amazon EC2 worker nodes.

Correct Answer: *AB*

Community vote distribution

AB (69%) AC (31%)

Comments

xBUGx Highly Voted 1 year, 2 months ago

Selected Answer: AC

I don't want confuse other...

upvoted 10 times

JohnYu 7 months, 3 weeks ago

AWS Lambda can run Docker containers, but it is more suited for short-duration, event-driven functions rather than long-running services.
While you can use Lambda to run container images, it is typically used for microservices or functions rather than managing scalable containerized applications in the same way that ECS or EKS does.

upvoted 4 times

FlyingHawk Most Recent 5 months, 1 week ago

Selected Answer: AB

A or B is okay, why not C?
While AWS Lambda now supports running Linux-based container images, it is primarily designed for microservices or event-driven functions rather than managing scalable containerized applications like Amazon ECS or Amazon EKS
upvoted 2 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: AB

A or B is OK.
About C: Amazon API Gateway and AWS Lambda are not typically used for running Docker containers. Lambda is designed for serverless functions, and while you can use Lambda to invoke containerized tasks, it does not run containers directly.
upvoted 2 times

JA2018 6 months, 1 week ago

Selected Answer: AB

AWS Fargate:
Both ECS and EKS with Fargate provide a fully managed container orchestration service where you don't need to manage underlying infrastructure like EC2 instances, making it ideal for scaling and reducing operational overhead.
Scaling capabilities:

Both options allow for automatic scaling based on demand, meaning the container instances will automatically scale up or down based on application requirements.
upvoted 1 times

striker89 7 months, 1 week ago

Selected Answer: AB

API Gateway + Lamda are not designed for running Dockers containers directly
upvoted 3 times

MatAlves 8 months, 3 weeks ago

Selected Answer: AC

- Container = ECS or EKS
- "managed service to host the application" = Fargate
- "not result in additional operational overhead or infrastructure to manage" = ECS for the win.

The main difference between ECS and EKS = simplicity vs flexibility.

https://aws.amazon.com/blogs/containers/amazon-ecs-vs-amazon-eks-making-sense-of-aws-container-services/
upvoted 1 times

MatAlves 8 months, 3 weeks ago

now, if instead of asking the "best" (combination) solution you want TWO SOLUTIONS, then I see value in A-B.
upvoted 2 times

Salilgen 5 months, 1 week ago

Where do you read "best" solution?
upvoted 1 times

Abdullah2004 9 months, 3 weeks ago

Selected Answer: AB

For sure
upvoted 3 times

MatAlves 8 months, 3 weeks ago

Why so many trolls recently...
upvoted 1 times

Abdullah2004 7 months, 2 weeks ago

Dear MatAlves before you describe my answer as trolls try to study and focus
I will explain to you that C is not correct
Lambda is for running code in response to events and does not natively support running Docker containers in the traditional sense of a containerized application. Lambda can use container image as deployment packages but this setup isn't ideal for managing and calling complex containerized applications
So it's A + B

upvoted 4 times

1e22522 10 months ago

Selected Answer: AB

glowies out reeeeeeeeeeeeeeee

upvoted 2 times

Rhydian25 11 months, 1 week ago

Selected Answer: AB

The question is asking for two alternatives to run Docker containers in a serverless service with minimal effort.

Option C will require a lot of effort to configure the Lambda and the API Gateway to run the Container correctly.

Instead, just use EKS o ECS with Fargate to execute the container image

upvoted 4 times

victor78 12 months ago

It should be AB

upvoted 2 times

sheilawu 12 months ago

Selected Answer: AB

It should be AB, Container :ECS , EKS

upvoted 2 times

sheilawu 12 months ago

No, I decided to change my option since this question is asking "Docker"&"AWS"
it is not asking local Kubernet, so AC should be the right answer.

upvoted 1 times

jcck202020 1 year, 1 month ago

AB are using AWS Fargate which IS considered a managed service, option C does not run containers, , DE you have to manage your own EC2 instances thus not consider managed

upvoted 2 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: AB

Everyone is picking AB too..

upvoted 2 times

[Removed] 1 year, 3 months ago

Selected Answer: AB

Option AB

upvoted 3 times

NayeraB 1 year, 3 months ago

Are people picking A&B as alternate solutions? Is the question asking for alternates?? Am I missing something? Somebody explain please I'm super confused.

upvoted 2 times

Drew3000 1 year, 2 months ago

I believe so. Based on other questions, they would have asked "which combination"

I believe so. Based on other questions, they would have asked "Which combination

upvoted 2 times

Cali182 1 year, 3 months ago

The question states itself. Which Solutions....?

upvoted 3 times

kempes 1 year, 4 months ago

Option AB

upvoted 2 times

Andy_09 1 year, 4 months ago

Option AB

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #773 *Topic 1*

An ecommerce company is running a seasonal online sale. The company hosts its website on Amazon EC2 instances spanning multiple Availability Zones. The company wants its website to manage sudden traffic increases during the sale.

Which solution will meet these requirements MOST cost-effectively?

A. Create an Auto Scaling group that is large enough to handle peak traffic load. Stop half of the Amazon EC2 instances. Configure the Auto Scaling group to use the stopped instances to scale out when traffic increases.

B. Create an Auto Scaling group for the website. Set the minimum size of the Auto Scaling group so that it can handle high traffic volumes without the need to scale out.

C. Use Amazon CloudFront and Amazon ElastiCache to cache dynamic content with an Auto Scaling group set as the origin. Configure the Auto Scaling group with the instances necessary to populate CloudFront and ElastiCache. Scale in after the cache is fully populated.

D. Configure an Auto Scaling group to scale out as traffic increases. Create a launch template to start new instances from a preconfigured Amazon Machine Image (AMI). Most Voted

Correct Answer: *D*

Community vote distribution

D (89%) C (11%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option D

upvoted 8 times

[Removed] Highly Voted 1 year ago

Selected Answer: D

The most cost-effective solution is:

D. Configure an Auto Scaling group to scale out as traffic increases. Create a launch template to start new instances from a preconfigured Amazon Machine Image (AMI).

upvoted 6 times

JA2018 Most Recent 6 months, 1 week ago

Selected Answer: D

Reasons?

Cost-effective scaling: Allows the company to only pay for the EC2 instances needed during peak traffic periods, automatically scaling up when necessary and scaling down when traffic subsides, which is the most cost-efficient approach.

Launch template and AMI: Using a launch template with a preconfigured AMI simplifies the scaling process by ensuring new instances are quickly provisioned with the required configurations.

upvoted 2 times

viejito 1 year ago

En la respuesta C : Al usar servicios como Amazon CloudFront y Amazon ElastiCache para almacenar en caché el contenido dinámico, reduciendo la carga en las instancias de Amazon EC2 y mejorando la velocidad de entrega del contenido a los usuarios finales. Esto resulta en una solución más rentable y eficiente en comparación con la respuesta D : simplemente escalar instancias de EC2 sin considerar medidas adicionales para optimizar el rendimiento y reducir los costos . Por lo que la opción correcta es la C .

upvoted 1 times

sandordini 1 year, 1 month ago

Selected Answer: D

Cloudfront coulld be a good idea, but it seems to be a simple scaling scenario.. IMO: D

upvoted 3 times

aditianand 1 year ago

Don't we need a launch template? How can it just launch from AMI image

upvoted 2 times

buzzinmumbai 1 year, 1 month ago

Answer is D .C is not cost effective to use elasticache .Not sure if you can have ASG as the origin.

upvoted 2 times

geraltRebo 1 year, 1 month ago

Selected Answer: D

Sorry D

upvoted 2 times

geraltRebo 1 year, 1 month ago

Selected Answer: C

Option C

upvoted 1 times

TruthWS 1 year, 2 months ago

Option D bring a most cost effective

upvoted 2 times

JCAWS 1 year, 2 months ago

Selected Answer: C

C more suitable

upvoted 1 times

stephensimudemy 1 year, 3 months ago

Selected Answer: D

It's D

upvoted 3 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #774 *Topic 1*

A solutions architect must provide an automated solution for a company's compliance policy that states security groups cannot include a rule that allows SSH from 0.0.0.0/0. The company needs to be notified if there is any breach in the policy. A solution is needed as soon as possible.

What should the solutions architect do to meet these requirements with the LEAST operational overhead?

A. Write an AWS Lambda script that monitors security groups for SSH being open to 0.0.0.0/0 addresses and creates a notification every time it finds one.

B. Enable the restricted-ssh AWS Config managed rule and generate an Amazon Simple Notification Service (Amazon SNS) notification when a noncompliant rule is created. **Most Voted**

C. Create an IAM role with permissions to globally open security groups and network ACLs. Create an Amazon Simple Notification Service (Amazon SNS) topic to generate a notification every time the role is assumed by a user.

D. Configure a service control policy (SCP) that prevents non-administrative users from creating or editing security groups. Create a notification in the ticketing system when a user requests a rule that needs administrator permissions.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

Andy_09 Highly Voted 10 months, 1 week ago

Option B

upvoted 8 times

kempes Highly Voted 10 months ago

Selected Answer: B

Option B

upvoted 5 times

sandordini Most Recent 7 months, 2 weeks ago

Selected Answer: B

The others sound 'silly'... to say the least

upvoted 2 times

asdfcdsxdfc 9 months ago

Selected Answer: B

B looks correct

upvoted 2 times

Naveena_Devanga 9 months, 2 weeks ago

Option B
https://docs.aws.amazon.com/config/latest/developerguide/restricted-ssh.html

upvoted 4 times

hajra313 10 months ago

option b

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #775 *Topic 1*

Use Amazon Elastic Kubernetes Service (Amazon EKS) with Amazon EC2 worker nodes.

A company has deployed an application in an AWS account. The application consists of microservices that run on AWS Lambda and Amazon Elastic Kubernetes Service (Amazon EKS). A separate team supports each microservice. The company has multiple AWS accounts and wants to give each team its own account for its microservices.

A solutions architect needs to design a solution that will provide service-to-service communication over HTTPS (port 443). The solution also must provide a service registry for service discovery.

Which solution will meet these requirements with the LEAST administrative overhead?

A. Create an inspection VPC. Deploy an AWS Network Firewall firewall to the inspection VPC. Attach the inspection VPC to a new transit gateway. Route VPC-to-VPC traffic to the inspection VPC. Apply firewall rules to allow only HTTPS communication.

B. Create a VPC Lattice service network. Associate the microservices with the service network. Define HTTPS listeners for each service. Register microservice compute resources as targets. Identify VPCs that need to communicate with the services. Associate those VPCs with the service network. **Most Voted**

C. Create a Network Load Balancer (NLB) with an HTTPS listener and target groups for each microservice. Create an AWS PrivateLink endpoint service for each microservice. Create an interface VPC endpoint in each VPC that needs to consume that microservice.

D. Create peering connections between VPCs that contain microservices. Create a prefix list for each service that requires a connection to a client. Create route tables to route traffic to the appropriate VPC. Create security groups to allow only HTTPS communication.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

1dd Highly Voted 9 months ago

Selected Answer: B