VPC Lattice is a completely new way to simplify API communication between services or microservices in one or more AWS accounts.

upvoted 7 times

zinabu Most Recent 7 months, 2 weeks ago

Selected Answer: B

Amazon VPC Lattice is a new capability of Amazon Virtual Private Cloud (Amazon VPC) designed to simplify networking for service-to-service communication.
link: https://www.bing.com/search?q=what+VPC+Lattice+service+used+for+microservices&cvid=d706d95737274f388660cbda9b7b2c4e&gs_lcrp=EgZjaHJvbWUyBggAEEUYOTIICAEQ6QcY_FXSAQkyMTY1N2owajSoAgCwAgE&FORM=ANAB01&PC=U531

upvoted 4 times

aditianand 6 months, 3 weeks ago

Did you complete the exam recently? Was examtopics useful?

upvoted 2 times

phoenix2023 6 months, 3 weeks ago

Please keep in mind this is for helpful answers to THIS specific question. Please don't abuse it as a random info pitch for yourself. This is distracting and wasting others' time. Please respect other people's time.

upvoted 7 times

stephensimudemy 9 months, 3 weeks ago

Selected Answer: B

IT's B. Google VPC Lattice service network

upvoted 3 times

Andy_09 10 months, 1 week ago

Option B

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #776 *Topic 1*

A company has a mobile game that reads most of its metadata from an Amazon RDS DB instance. As the game increased in popularity, developers noticed slowdowns related to the game's metadata load times. Performance metrics indicate that simply scaling the database will not help. A solutions architect must explore all options that include capabilities for snapshots, replication, and sub-millisecond response times.

What should the solutions architect recommend to solve these issues?

A. Migrate the database to Amazon Aurora with Aurora Replicas.

B. Migrate the database to Amazon DynamoDB with global tables.

C. Add an Amazon ElastiCache for Redis layer in front of the database. **Most Voted**

D. Add an Amazon ElastiCache for Memcached layer in front of the database.

Correct Answer: *C*

Community vote distribution

C (80%) B (20%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option C is better as we need replication and snapshots

upvoted 20 times

JA2018 6 months, 1 week ago

Not sure if this is applicable to the stem. While Redis can be used as a cache, placing it in front of a relational database like RDS might not be enough to address the performance issues if the primary bottleneck is database access time.

upvoted 1 times

arunkpskpm 1 year, 3 months ago

C is correct as only Redis support snapshot feature :https://aws.amazon.com/elasticache/redis-vs-memcached/

upvoted 6 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: C

A - If the root cause is database-related (for example, read-heavy traffic), then OK. However, "Performance metrics indicate that simply scaling the database will not help." means that maybe caching is a more effective option.
B - Migrating from a relational database to DynamoDB requires so much effort. Such a headache.
C - ElastiCache for Redis supports replication and snapshots, and of course, sub-millisecond response times.
D - Memcached lacks built-in persistence, which is why game applications rarely use it.

upvoted 2 times

tch 5 months, 2 weeks ago

Selected Answer: C

capabilities for snapshots, replication

upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: B

Why I choose B?

1 High-performance, low latency: DynamoDB is a NoSQL database designed for fast, sub-millisecond response times, which is crucial for handling the high volume of metadata requests from a popular mobile game.

2 Scalability: DynamoDB scales automatically to handle fluctuating traffic without the need to manually manage database instances, making it ideal for a rapidly growing game.

3 Global tables: This feature allows for geographically distributed data access, ensuring quick response times for players across different regions.

upvoted 2 times

JunsK1e 11 months ago

Selected Answer: C

C is correct

upvoted 1 times

f07ed8f 1 year ago

Thinking the capabilities with "snapshots, replication, and sub-millisecond response times" is for the Database or selected solution(ElastiCache).

upvoted 2 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: C

C is correct

upvoted 2 times

nbellaiche 1 year, 3 months ago

Selected Answer: C

Réponse C

upvoted 1 times

osmk 1 year, 3 months ago

Selected Answer: C

:https://aws.amazon.com/elasticache/redis-vs-memcached/

upvoted 1 times

Andy_09 1 year, 4 months ago

Option D

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #777 *Topic 1*

A company uses AWS Organizations for its multi-account AWS setup. The security organizational unit (OU) of the company needs to share approved Amazon Machine Images (AMIs) with the development OU. The AMIs are created by using AWS Key Management Service (AWS KMS) encrypted snapshots.

Which solution will meet these requirements? (Choose two.)

A. Add the development team's OU Amazon Resource Name (ARN) to the launch permission list for the AMIs. **Most Voted**

B. Add the Organizations root Amazon Resource Name (ARN) to the launch permission list for the AMIs.

C. Update the key policy to allow the development team's OU to use the AWS KMS keys that are used to decrypt the snapshots. **Most Voted**

D. Add the development team's account Amazon Resource Name (ARN) to the launch permission list for the AMIs.

E. Recreate the AWS KMS key. Add a key policy to allow the Organizations root Amazon Resource Name (ARN) to use the AWS KMS key.

Correct Answer: *AC*

Community vote distribution

AC (100%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Changing to options AC

upvoted 13 times

osmk Highly Voted 1 year, 3 months ago

Selected Answer: AC

c=>https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/share-amis-with-organizations-and-OUs.html#allow-org-ou-to-use-key

A-->https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/share-amis-with-organizations-and-OUs.html#share-amis-org-ou

upvoted 5 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: AC

A - The security OU is explicitly allowing the development OU to use the AMI.
B - OK but too broad, and not secure.
C - Now they can use the encrypted snapshots.
D - Would be too much trouble if there are many accounts. A is enough.
E - Recreate the AWS KMS key? Uhm, why?

upvoted 2 times

Scheldon 11 months, 3 weeks ago

Selected Answer: AC

AnswerAC

Option A will allow to run/lunch AMIs
Option C will allow to decript AMIs which is necessery to run AMI.

upvoted 2 times

cjace 11 months, 4 weeks ago

CD - Solution C: Update the Key Policy
Why: The AMIs are created using KMS-encrypted snapshots, so the KMS keys must allow the development team's accounts to use these keys for decrypting the snapshots.
How: Update the key policy of the KMS key to include permissions for the development OU or specific accounts within that OU. This will enable those accounts to use the KMS key for decrypting the snapshots associated with the AMIs.
Solution D: Add the Development Team’s Account ARN to the Launch Permission List
Why: To share the AMIs with the development accounts, you need to grant launch permissions to those accounts. This allows the specified accounts to use the shared AMIs to launch instances.
How: Add the ARNs of the development team's accounts to the launch permission list of the AMIs. This can be done using the modify-image-attribute command in the AWS CLI, specifying the account IDs that should have launch permissions.

upvoted 1 times

Mikado211 1 year, 2 months ago

Selected Answer: AC

A : give users the right to launch
C : give users the right to decrypt

upvoted 4 times

Andy_09 1 year, 4 months ago

Option CD

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #778 *Topic 1*

A data analytics company has 80 offices that are distributed globally. Each office hosts 1 PB of data and has between 1 and 2 Gbps of internet bandwidth.

The company needs to perform a one-time migration of a large amount of data from its offices to Amazon S3. The company must complete the migration within 4 weeks.

Which solution will meet these requirements MOST cost-effectively?

A. Establish a new 10 Gbps AWS Direct Connect connection to each office. Transfer the data to Amazon S3.

B. Use multiple AWS Snowball Edge storage-optimized devices to store and transfer the data to Amazon S3. **Most Voted**

C. Use an AWS Snowmobile to store and transfer the data to Amazon S3.

D. Set up an AWS Storage Gateway Volume Gateway to transfer the data to Amazon S3.

Correct Answer: *B*

Community vote distribution

B (91%) C (9%)

Comments

mestule Highly Voted 1 year, 4 months ago

Selected Answer: B

B because too many offices that are geographically separated.

"data analytics company has 80 offices that are distributed globally."

upvoted 16 times

Andy_09 1 year, 4 months ago

Nice spot...completely missed that part !!

upvoted 4 times

MatAlves Most Recent 8 months, 3 weeks ago

Selected Answer: B

AWS Snowmobile: ideal for migrating big datasets containing 10PB or more and stored in one location. Snowmobile can help

AWS Snowmobile: ideal for migrating big datasets containing 10PB or more and stored in one location. Snowmobile can help you migrate all of these large datasets at once, but the process requires a high-speed backbone with hundreds of Gb/s of spare throughput.

AWS Snowball: ideal for datasets storing less than 10PB or datasets distributed across multiple locations. You can use Snowball to migrate data incrementally—this is a good alternative if you do not have enough bandwidth on the network backbone.

https://bluexp.netapp.com/blog/aws-cvo-blg-aws-snowball-vs-snowmobile-data-migration-options-comparedwork backbone.

upvoted 3 times

ccceb01 9 months, 2 weeks ago

Selected Answer: B

10 Snowball price (100TB x 10) still cheaper then 1 Snowmobile price ($4,100 plus additional cost for data transfer)

upvoted 2 times

aditianand 1 year ago

Why not D? Why not AWS storage gateway? With 1Gbps, they can transfer 1.25 GBPS which translates to 2.8 PB in 4 weeks. They need just 1 PB to be transferred

upvoted 2 times

buzzinmumbai 1 year, 1 month ago

As of March 2024 AWS has stopped offering snowmobile as a service .So B is the right answer.Hopefully they don't ask this question :)

upvoted 3 times

Tanidanindo 1 year, 1 month ago

Selected Answer: C

Too large for snowball devices.

upvoted 2 times

Naveena_Devanga 1 year, 3 months ago

Option C,
An AWS Snowmobile has a maximum storage capacity of 100 petabytes (PB). This is equivalent to the capacity of 1,250 Snowball Edge devices

upvoted 1 times

HarryLopez 1 year, 3 months ago

but there are many offices geographically distributed, so snowmobile for each one of them adds up to a lot of cost as compared to option B).

upvoted 3 times

sandordini 1 year, 1 month ago

Snowmobile advised over 10PB! Definitely snowball

upvoted 1 times

chefKC 1 year, 4 months ago

option B

upvoted 2 times

Andy_09 1 year, 4 months ago

Option C looks good, as option B would lead to usage of too many snowball devices.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #779 *Topic 1*

A company has an Amazon Elastic File System (Amazon EFS) file system that contains a reference dataset. The company has applications on Amazon EC2 instances that need to read the dataset. However, the applications must not be able to change the dataset. The company wants to use IAM access control to prevent the applications from being able to modify or delete the dataset.

Which solution will meet these requirements?

A. Mount the EFS file system in read-only mode from within the EC2 instances.

B. Create a resource policy for the EFS file system that denies the elasticfilesystem:ClientWrite action to the IAM roles that are attached to the EC2 instances. **Most Voted**

C. Create an identity policy for the EFS file system that denies the elasticfilesystem:ClientWrite action on the EFS file system.

D. Create an EFS access point for each application. Use Portable Operating System Interface (POSIX) file permissions to allow read-only access to files in the root directory.

Correct Answer: *B*

Community vote distribution

B (73%) C (23%) D (5%)

Comments

lenotc Highly Voted 1 year, 2 months ago

Selected Answer: B

B correct best solution best well architected
C wrong because identity policies are typically associated with users or roles, not directly with the EFS file system
D wrong because POSIX file permissions at the root directory level may not be sufficient to prevent modifications to other directories or files
A is so far away

upvoted 6 times

hajra313 Highly Voted 1 year, 4 months ago

Create an EFS access point for each application. Use Portable Operating System Interface (POSIX) file permissions to allow read-only access to files in the root directory.

Explanation:
By creating an EFS access point for each application and configuring POSIX file permissions to allow read-only access, you can enforce the desired access control. This approach restricts write and delete actions on the dataset while allowing read access, aligning with the company's requirements.
upvoted 6 times

MatAlves 8 months, 3 weeks ago

Resource policies are included in " IAM to control":

"Using IAM to control file system data access

NFS clients can identify themselves using an IAM role when connecting to an EFS file system. When a client connects to a file system, Amazon EFS evaluates the file system’s IAM resource policy, which is called a file system policy, along with any identity-based IAM policies to determine the appropriate file system access permissions to grant."

https://docs.aws.amazon.com/efs/latest/ug/iam-access-control-nfs-efs.html
upvoted 2 times

f07ed8f 1 year ago

Please note that the question is asking "The company wants to use IAM access control to prevent the applications from being able to modify or delete the dataset."
upvoted 5 times

MatAlves Most Recent 8 months, 3 weeks ago

Selected Answer: B

- Identity-based policies are attached to an IAM user, group, or role.
- Resource-based policies are attached to a resource.
- elasticfilesystem:ClientWrite: Provides write permissions on a file system.

EFS is a RESOURCE, so that excludes "C" (we need a resource policy).

https://docs.aws.amazon.com/efs/latest/ug/iam-access-control-nfs-efs.html

https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_identity-vs-resource.html
upvoted 2 times

elmyth 9 months ago

Selected Answer: B

There is no such thing as an "identity policy" for EFS.
upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: B

2 ways to prevent writing to the file system:
1. The mount option in the /etc/fstab file is set to read-only access. > A
2. IAM policy indicates read-only access, or root access disabled. > B
The question clearly states they are looking to use IAM access control
upvoted 3 times

Ansuman_lucky 1 year, 2 months ago

prevent the applications from being able to modify or delete the dataset.-- This means a role would be used. So answer is B
upvoted 4 times

xBUGx 1 year, 2 months ago

IAM policies are used to control access to AWS resources, including Amazon EFS. By default, IAM policies control access to the EFS API actions, such as elasticfilesystem:ClientWrite, which allows clients to write to the file system. However, POSIX file permissions control access to files within the file system itself, which is independent of IAM policies.

While using POSIX file permissions can restrict access to the files within the file system, it doesn't prevent a user or application with the appropriate IAM permissions from modifying or deleting those files directly through the EFS API.
upvoted 4 times

HarryLopez 1 year, 3 months ago

Selected Answer: B

B)
IAM needs to be used, so A) & D) are out.
So b/w B) and C), Resource policies are meant for specific aws service or resource while Identity policies are attached to an identity (user, group or role). C) attached identity policy to EFS, dont know how and why. Hence, B).

upvoted 3 times

osmk 1 year, 3 months ago

Selected Answer: C

company wants to use IAM access control to prevent https://docs.aws.amazon.com/efs/latest/ug/iam-access-control-nfs-efs.html

upvoted 3 times

jaswantn 1 year, 3 months ago

Selected Answer: D

option D

upvoted 1 times

Salilgen 5 months, 1 week ago

You use IAM policy to control permissions with access points
https://docs.aws.amazon.com/efs/latest/ug/access-points-iam-policy.html

upvoted 1 times

Oo_Cc 1 year, 3 months ago

Selected Answer: C

"The company wasn't to use IAM access control". Yes, it would deny writing action to everything .. but it's still the only one that uses IAM.

upvoted 2 times

MatAlves 8 months, 3 weeks ago

"Using IAM to control file system data access

NFS clients can identify themselves using an IAM role when connecting to an EFS file system. When a client connects to a file system, Amazon EFS evaluates the file system's IAM resource policy, which is called a file system policy, along with any identity-based IAM policies to determine the appropriate file system access permissions to grant."

https://docs.aws.amazon.com/efs/latest/ug/iam-access-control-nfs-efs.html

upvoted 2 times

MatAlves 8 months, 3 weeks ago

What we need to change is the " IAM resource policy".

upvoted 1 times

Andy_09 1 year, 4 months ago

Option B

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #780 *Topic 1*

A company has hired an external vendor to perform work in the company's AWS account. The vendor uses an automated tool that is hosted in an AWS account that the vendor owns. The vendor does not have IAM access to the company's AWS account. The company needs to grant the vendor access to the company's AWS account.

Which solution will meet these requirements MOST securely?

A. Create an IAM role in the company's account to delegate access to the vendor's IAM role. Attach the appropriate IAM policies to the role for the permissions that the vendor requires. **Most Voted**

B. Create an IAM user in the company's account with a password that meets the password complexity requirements. Attach the appropriate IAM policies to the user for the permissions that the vendor requires.

C. Create an IAM group in the company's account. Add the automated tool's IAM user from the vendor account to the group. Attach the appropriate IAM policies to the group for the permissions that the vendor requires.

D. Create an IAM user in the company's account that has a permission boundary that allows the vendor's account. Attach the appropriate IAM policies to the user for the permissions that the vendor requires.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

Mikado211 Highly Voted 1 year, 1 month ago

Selected Answer: A

When you have somebody from another account who needs a resource in your account
- create a role to access to this account
- allow the remote account to asume the role.

upvoted 8 times

Andy_09 Highly Voted 1 year, 4 months ago

Option A looks ok

upvoted 7 times

Salilgen Most Recent 5 months, 1 week ago

Selected Answer: A

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_common-scenarios_third-party.html

upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: A

Least privilege principle: Using an IAM role instead of a user adheres to the least privilege principle as it only grants the vendor access to specific actions needed for their task, not full account access like a user would have.

No direct access to the company account: By delegating access through a role, the vendor's IAM user in their own account does not need to have direct login credentials to the company's AWS account, enhancing security.

Granular control: IAM policies attached to the role can be carefully crafted to provide only the necessary permissions for the vendor's automated tool, limiting potential damage in case of unauthorized access

upvoted 1 times

Scheldon 11 months, 3 weeks ago

Selected Answer: A

AnswerA

I would go with option A

upvoted 2 times

osmk 1 year, 3 months ago

Selected Answer: A

Question #222

upvoted 6 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #781 *Topic 1*

A company wants to run its experimental workloads in the AWS Cloud. The company has a budget for cloud spending. The company's CFO is concerned about cloud spending accountability for each department. The CFO wants to receive notification when the spending threshold reaches 60% of the budget.

Which solution will meet these requirements?

A. Use cost allocation tags on AWS resources to label owners. Create usage budgets in AWS Budgets. Add an alert threshold to receive notification when spending exceeds 60% of the budget. **Most Voted**

B. Use AWS Cost Explorer forecasts to determine resource owners. Use AWS Cost Anomaly Detection to create alert threshold notifications when spending exceeds 60% of the budget.

C. Use cost allocation tags on AWS resources to label owners. Use AWS Support API on AWS Trusted Advisor to create alert threshold notifications when spending exceeds 60% of the budget.

D. Use AWS Cost Explorer forecasts to determine resource owners. Create usage budgets in AWS Budgets. Add an alert threshold to receive notification when spending exceeds 60% of the budget.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

NayeraB Highly Voted 1 year, 3 months ago

Selected Answer: A

Nothing with cost explorer in it, and I don't want to be Captain Obvious but we need to set the budget alerts through AWS Budgets, so A

upvoted 8 times

Scheldon Most Recent 12 months ago

Selected Answer: A

AnswerA

https://docs.aws.amazon.com/cost-management/latest/userguide/budgets-controls.html

upvoted 3 times

Andy_09 1 year, 4 months ago

Option A

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #782 *Topic 1*

A company wants to deploy an internal web application on AWS. The web application must be accessible only from the company's office. The company needs to download security patches for the web application from the internet.

The company has created a VPC and has configured an AWS Site-to-Site VPN connection to the company's office. A solutions architect must design a secure architecture for the web application.

Which solution will meet these requirements?

A. Deploy the web application on Amazon EC2 instances in public subnets behind a public Application Load Balancer (ALB). Attach an internet gateway to the VPC. Set the inbound source of the ALB's security group to 0.0.0.0/0.

B. Deploy the web application on Amazon EC2 instances in private subnets behind an internal Application Load Balancer (ALB). Deploy NAT gateways in public subnets. Attach an internet gateway to the VPC. Set the inbound source of the ALB's security group to the company's office network CIDR block. **Most Voted**

C. Deploy the web application on Amazon EC2 instances in public subnets behind an internal Application Load Balancer (ALB). Deploy NAT gateways in private subnets. Attach an internet gateway to the VPSet the outbound destination of the ALB's security group to the company's office network CIDR block.

D. Deploy the web application on Amazon EC2 instances in private subnets behind a public Application Load Balancer (ALB). Attach an internet gateway to the VPC. Set the outbound destination of the ALB's security group to 0.0.0.0/0.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option B

upvoted 7 times

osmk Highly Voted 1 year, 3 months ago

Selected Answer: B

none sense why IGW on top of NATGW.

upvoted 6 times

MatAlves 8 months, 3 weeks ago

https://docs.aws.amazon.com/network-firewall/latest/developerguide/arch-igw-ngw.html

Confusing, I agree. But it seems to be recommended in some cases.
upvoted 1 times

striker89 Most Recent 7 months, 1 week ago

Selected Answer: B

I Would go for B even if NAT GW allow outbound traffic ONLY. Still wondering how the Company newtwork will access Private Subnet in the VPC.
upvoted 1 times

Scheldon 11 months, 3 weeks ago

AnswerB

Server and LB in Private will hide WEB application from the word. NAT will allow for server's access to the internet in case of need
upvoted 1 times

NayeraB 1 year, 3 months ago

Selected Answer: B

B is well structured
upvoted 3 times

ogerber 1 year, 3 months ago

To my opinion, with only having inbound of the companys CIDR block, it will not include access for the patches available online.
i would go for D
upvoted 4 times

sandordini 1 year, 1 month ago

Incorrect: B says inbound, D says outbound. Outbound for ALB are the EC2 Instances.
upvoted 1 times

kempes 1 year, 4 months ago

Selected Answer: B

Option B
upvoted 4 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #783 *Topic 1*

A company maintains its accounting records in a custom application that runs on Amazon EC2 instances. The company needs to migrate the data to an AWS managed service for development and maintenance of the application data. The solution must require minimal operational support and provide immutable, cryptographically verifiable logs of data changes.

Which solution will meet these requirements MOST cost-effectively?

A. Copy the records from the application into an Amazon Redshift cluster.

B. Copy the records from the application into an Amazon Neptune cluster.

C. Copy the records from the application into an Amazon Timestream database.

D. Copy the records from the application into an Amazon Quantum Ledger Database (Amazon QLDB) ledger. Most Voted

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option D

upvoted 7 times

asdfcdsxdfc Highly Voted 1 year, 3 months ago

Selected Answer: D

Amazon QLDB
• QLDB stands for "Quantum Ledger Database"
• A ledger is a book recording financial transactions
• Fully Managed, Serverless, High available, Replication across 3 AZ
• Used to review history of all the changes made to your application data over time
• Immutable system: no entry can be removed or modified, cryptographically verifiable

upvoted 5 times

Scheldon Most Recent 11 months, 3 weeks ago

Selected Answer: D

AnswerD

Amazon Quantum Ledger Database (Amazon QLDB) is a fully managed ledger database that provides a transparent, immutable, and cryptographically verifiable transaction log owned by a central trusted authority.

https://docs.aws.amazon.com/qldb/latest/developerguide/what-is.html

upvoted 4 times

f07ed8f 1 year ago

Selected Answer: D

Amazon Quantum Ledger Database (Amazon QLDB) is a fully managed ledger database that provides a transparent, immutable, and cryptographically verifiable transaction log.

upvoted 3 times

Mikado211 1 year, 2 months ago

Selected Answer: D

immutable, cryptographically verifiable ==> Amazon QLDB

upvoted 3 times

[Removed] 1 year, 3 months ago

Selected Answer: D

https://aws.amazon.com/qldb/
Amazon Quantum Ledger Database (Amazon QLDB) is a fully managed ledger database that provides a transparent, immutable, and cryptographically verifiable transaction log.

upvoted 3 times

NayeraB 1 year, 3 months ago

Selected Answer: D

D is correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #784 *Topic 1*

A company's marketing data is uploaded from multiple sources to an Amazon S3 bucket. A series of data preparation jobs aggregate the data for reporting. The data preparation jobs need to run at regular intervals in parallel. A few jobs need to run in a specific order later.

The company wants to remove the operational overhead of job error handling, retry logic, and state management.

Which solution will meet these requirements?

A. Use an AWS Lambda function to process the data as soon as the data is uploaded to the S3 bucket. Invoke other Lambda functions at regularly scheduled intervals.

B. Use Amazon Athena to process the data. Use Amazon EventBridge Scheduler to invoke Athena on a regular internal.

C. Use AWS Glue DataBrew to process the data. Use an AWS Step Functions state machine to run the DataBrew data preparation jobs. Most Voted

D. Use AWS Data Pipeline to process the data. Schedule Data Pipeline to process the data once at midnight.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

[Removed] Highly Voted 1 year, 3 months ago

Selected Answer: C

data preparation = Glue DataBrew https://docs.aws.amazon.com/databrew/latest/dg/what-is.html
state handling = DataBrew with Step Functions https://docs.aws.amazon.com/step-functions/latest/dg/connect-databrew.html

upvoted 12 times

Andy_09 Highly Voted 1 year, 4 months ago

Option C

upvoted 9 times

FlyingHawk Most Recent 4 months, 1 week ago

Selected Answer: C

Both c and d are workable, but AWS Data Pipeline is considered outdated because AWS has introduced newer, more efficient, and serverless alternatives for data orchestration and ETL (Extract, Transform, Load) tasks. While AWS Data Pipeline is still available, it is not actively promoted or updated by AWS, and many companies have shifted to AWS Glue, Step Functions, and EventBridge for data workflow automation.

https://aws.amazon.com/blogs/big-data/build-event-driven-data-quality-pipelines-with-aws-glue-databrew/

upvoted 2 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: C

A - Lambda function alone is incompetent for retries or running in a specific order.
B - Athena cannot automatically handle errors or running in a specific order.
C - YES and fully managed. See below.
D - Data Pipeline with retries and error handling is A LOT OF WORK.
Takeaway:
1. Glue DataBrew simplifies the data preparation process with a visual interface, and it integrates with S3 for input/output.
2. Step Functions efficiently manages parallel and sequential workflows while minimizing operational overhead for state management, error handling, and retries.

upvoted 3 times

Scheldon 11 months, 3 weeks ago

Selected Answer: C

AnswerC

With Step Functions' built-in controls, you can examine the state of each step in your workflow to make sure that your application runs in order and as expected.

https://docs.aws.amazon.com/step-functions/latest/dg/welcome.html

AWS Glue is a serverless data integration service that makes it easy for analytics users to discover, prepare, move, and integrate data from multiple sources.

https://docs.aws.amazon.com/glue/latest/dg/what-is-glue.html

upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: C

c looks correct

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #785 *Topic 1*

A solutions architect is designing a payment processing application that runs on AWS Lambda in private subnets across multiple Availability Zones. The application uses multiple Lambda functions and processes millions of transactions each day.

The architecture must ensure that the application does not process duplicate payments.

Which solution will meet these requirements?

A. Use Lambda to retrieve all due payments. Publish the due payments to an Amazon S3 bucket. Configure the S3 bucket with an event notification to invoke another Lambda function to process the due payments.

B. Use Lambda to retrieve all due payments. Publish the due payments to an Amazon Simple Queue Service (Amazon SQS) queue. Configure another Lambda function to poll the SQS queue and to process the due payments.

C. Use Lambda to retrieve all due payments. Publish the due payments to an Amazon Simple Queue Service (Amazon SQS) FIFO queue. Configure another Lambda function to poll the FIFO queue and to process the due payments. Most Voted

D. Use Lambda to retrieve all due payments. Store the due payments in an Amazon DynamoDB table. Configure streams on the DynamoDB table to invoke another Lambda function to process the due payments.

Correct Answer: *C*

Community vote distribution

C (82%) D (18%)

Comments

hajra313 Highly Voted 1 year, 4 months ago

Standard queues provide at-least-once delivery, which means that each message is delivered at least once.

FIFO queues provide exactly-once processing , which means that each message is delivered once and remains available until a consumer processes it and deletes it. Duplicates are not introduced into the queue. OPTION C

upvoted 18 times

Scheldon Most Recent 11 months, 3 weeks ago

Selected Answer: C

AnswerC

SQS FIFO was created for such tasks

Unlike standard queues, FIFO queues don't introduce duplicate messages. FIFO queues help you avoid sending duplicates to a queue. If you retry the SendMessage action within the 5-minute deduplication interval, Amazon SQS doesn't introduce any duplicates into the queue.
https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/FIFO-queues-exactly-once-processing.html

upvoted 3 times

escalibran 1 year, 2 months ago

Selected Answer: C

C over D, because
https://docs.aws.amazon.com/lambda/latest/dg/with-ddb.html Processing dynamo streams with lambda can cause duplication.
SQS FIFO can be configured for High Throughput to exceed the 3000/s (batched) limit
https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/high-throughput-fifo.html

I previously worked with payments and would argue that either option doesn't fully solve duplications. Events might be sent multiple times from source, you definitely want to perform de-duplication and have some sort of idempotent processing for them, instead of just blindly processing each thing you're given.

upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: C

c is correct

upvoted 2 times

shahreh1 1 year, 3 months ago

Option C:
FIFO queues
Exactly-Once Processing – A message is delivered once and remains available until a consumer processes and deletes it. Duplicates aren't introduced into the queue.
First-In-First-Out Delivery – The order in which messages are sent and received is strictly preserved.

upvoted 2 times

FZA24 1 year, 3 months ago

Selected Answer: C

Option C Fifo

upvoted 3 times

Mikado211 1 year, 3 months ago

SQS can have duplicate messages in case of problems with the timeout window.

upvoted 2 times

haci 1 year, 3 months ago

Selected Answer: C

"The application does not process duplicate payments" is the key point, which leads us directly to SQS FIFO

upvoted 3 times

Cali182 1 year, 3 months ago

Selected Answer: D

Option D
DynamoDB Streams helps ensure the following:

Each stream record appears exactly once in the stream.

For each item that is modified in a DynamoDB table, the stream records appear in the same sequence as the actual modifications to the item.

DynamoDB Streams writes stream records in near-real time so that you can build applications that consume these streams and take action based on the contents.

upvoted 3 times

c3c7e62 1 month ago

It seems that DynamoDB Streams deduplication is not guaranteed -
"A Lambda consumer for a DynamoDB stream doesn't guarantee exactly once delivery and may lead to occasional duplicates. Make sure your Lambda function code is idempotent to prevent unexpected issues from arising because of duplicate processing."

https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Streams.Lambda.BestPracticesWithDynamoDB.html
upvoted 1 times

jaswantn 1 year, 3 months ago

Option D...If you need to handle millions of transactions each day, you might need to consider other approach instead of SQS FIFO. And amongst the given options, we have DynmamoDB that maintains order in the streams.
upvoted 1 times

NayeraB 1 year, 3 months ago

I'm not sure if the answer is DynamoDB as well, but answering your question, SQS Fifo can handle 300 messages/second without batching, 3,000 messages/second with batching. Assuming we're using the 300/sec option, with 86,400 seconds in a day, that gives you 25,920,000 messages, so in short, yes SQS can handle millions of requests each day.
Not to mention DynamoDB doesn't provide the exactly-once processing the SQS offer and clearly requested in the question. That's just my train of thought, I'm happy to be corrected.
upvoted 3 times

jaswantn 1 year, 3 months ago

Dynamodb streams with partition key can be used to implement exactly once processing. There are many options with dynamodb to check for already processed item, and can be filtered out so that they are processed only once.
upvoted 1 times

jaswantn 1 year, 3 months ago

This calculation limits the number of transactions to 25 million a day. What if there are transactions exceeding this limit? As question say .... millions of transactions a day; that could be 70,80 or 90 millions also. In that case how SQS FIFO would perform?
Happy to be corrected with more convincing facts
upvoted 1 times

kempes 1 year, 4 months ago

Option c
upvoted 2 times

Andy_09 1 year, 4 months ago

Option B
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #786 *Topic 1*

A company runs multiple workloads in its on-premises data center. The company's data center cannot scale fast enough to meet the company's expanding business needs. The company wants to collect usage and configuration data about the on-premises servers and workloads to plan a migration to AWS.

Which solution will meet these requirements?

A. Set the home AWS Region in AWS Migration Hub. Use AWS Systems Manager to collect data about the on-premises servers.

B. Set the home AWS Region in AWS Migration Hub. Use AWS Application Discovery Service to collect data about the on-premises servers. **Most Voted**

C. Use the AWS Schema Conversion Tool (AWS SCT) to create the relevant templates. Use AWS Trusted Advisor to collect data about the on-premises servers.

D. Use the AWS Schema Conversion Tool (AWS SCT) to create the relevant templates. Use AWS Database Migration Service (AWS DMS) to collect data about the on-premises servers.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

Kezuko Highly Voted 1 year, 2 months ago

Still the planning stage, C and D is out.

upvoted 5 times

Scheldon Most Recent 11 months, 3 weeks ago

Selected Answer: B

AnswerB
AWS Migration Hub delivers a guided end-to-end migration and modernization journey through discovery, assessment, planning, and execution.
https://aws.amazon.com/migration-hub/
AWS Application Discovery Service helps you plan your migration to the AWS cloud by collecting usage and configuration data about your on-premises servers and databases. Application Discovery Service is integrated with AWS Migration Hub and AWS Database Migration Service Fleet Advisor

Database Migration Service Fleet Advisor.
https://docs.aws.amazon.com/application-discovery/latest/userguide/what-is-appdiscovery.html
upvoted 4 times

lpergorta 1 year, 2 months ago

Option D
upvoted 1 times

lpergorta 1 year, 2 months ago

Sorry B
upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: B

B is correct
upvoted 2 times

[Removed] 1 year, 3 months ago

Selected Answer: B

AWS Application Discovery Service helps you plan your migration to the AWS cloud by collecting usage and configuration data about your on-premises servers and databases. https://docs.aws.amazon.com/application-discovery/latest/userguide/what-is-appdiscovery.html
upvoted 4 times

Andy_09 1 year, 4 months ago

Option B
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #787 *Topic 1*

A company has an organization in AWS Organizations that has all features enabled. The company requires that all API calls and logins in any existing or new AWS account must be audited. The company needs a managed solution to prevent additional work and to minimize costs. The company also needs to know when any AWS account is not compliant with the AWS Foundational Security Best Practices (FSBP) standard.

Which solution will meet these requirements with the LEAST operational overhead?

A. Deploy an AWS Control Tower environment in the Organizations management account. Enable AWS Security Hub and AWS Control Tower Account Factory in the environment. **Most Voted**

B. Deploy an AWS Control Tower environment in a dedicated Organizations member account. Enable AWS Security Hub and AWS Control Tower Account Factory in the environment.

C. Use AWS Managed Services (AMS) Accelerate to build a multi-account landing zone (MALZ). Submit an RFC to self-service provision Amazon GuardDuty in the MALZ.

D. Use AWS Managed Services (AMS) Accelerate to build a multi-account landing zone (MALZ). Submit an RFC to self-service provision AWS Security Hub in the MALZ.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

FlyingHawk 4 months ago

Selected Answer: A

This question was in my today's exam.

upvoted 3 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: A

A - Control Tower enforces centralized logging, monitoring, and account management. Security Hub monitors compliance with AWS Foundational Security Best Practices (FSBP). Account Factory simplifies the creation of new accounts with preconfigured settings.

B - Control Tower is intended to be deployed in the management account, not a member account, which deviates from AWS-

B - Control Tower is intended to be deployed in the management account, not a member account, which deviates from AWS recommended practices.
C - GuardDuty focuses on security threats rather than compliance monitoring for FSBP.
D - AMS Accelerate is more suitable for enterprises with highly complex requirements. Also, more operational overheads and cost more, but the question wants "to prevent additional work and to minimize costs".
upvoted 1 times

LuongTo 7 months, 2 weeks ago

Why not B?
upvoted 1 times

Kezuko 1 year, 2 months ago

Selected Answer: A

https://docs.aws.amazon.com/controltower/latest/userguide/security-hub-controls.html
upvoted 4 times

lpergorta 1 year, 2 months ago

Option D
upvoted 1 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: A

A is correct
upvoted 2 times

kempes 1 year, 4 months ago

Selected Answer: A

Option A
upvoted 4 times

Andy_09 1 year, 4 months ago

Option A
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #788 *Topic 1*

A company has stored 10 TB of log files in Apache Parquet format in an Amazon S3 bucket. The company occasionally needs to use SQL to analyze the log files.

Which solution will meet these requirements MOST cost-effectively?

A. Create an Amazon Aurora MySQL database. Migrate the data from the S3 bucket into Aurora by using AWS Database Migration Service (AWS DMS). Issue SQL statements to the Aurora database.

B. Create an Amazon Redshift cluster. Use Redshift Spectrum to run SQL statements directly on the data in the S3 bucket.

C. Create an AWS Glue crawler to store and retrieve table metadata from the S3 bucket. Use Amazon Athena to run SQL statements directly on the data in the S3 bucket. **Most Voted**

D. Create an Amazon EMR cluster. Use Apache Spark SQL to run SQL statements directly on the data in the S3 bucket.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: C

A - Aurora is cool but migrating 10 TB of data incurs significant costs and operational overhead.
B - Redshift Spectrum allows querying data directly in S3 without loading it into Redshift, but costs are really high especially for infrequent use.
C - Athena is serverless and charges only for the data scanned by queries. Glue Crawler automatically extracts metadata and schema information from the Parquet files. No need to migrate anything.
D - Just by the look of it I know I'll go bankrupt if I choose that.

upvoted 1 times

sandordini 7 months, 2 weeks ago

Selected Answer: C

S3 + SQL = Athena

upvoted 3 times

Kezuko 8 months, 3 weeks ago

Selected Answer: C

Apache Parquet => Glue Crawler

upvoted 4 times

asdfcdsxdfc 9 months ago

Selected Answer: C

c is correct

upvoted 2 times

kempes 10 months ago

Selected Answer: C

Option C

upvoted 4 times

Andy_09 10 months, 1 week ago

Option C

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #789 *Topic 1*

A company needs a solution to prevent AWS CloudFormation stacks from deploying AWS Identity and Access Management (IAM) resources that include an inline policy or "*" in the statement. The solution must also prohibit deployment of Amazon EC2 instances with public IP addresses. The company has AWS Control Tower enabled in its organization in AWS Organizations.

Which solution will meet these requirements?

A. Use AWS Control Tower proactive controls to block deployment of EC2 instances with public IP addresses and inline policies with elevated access or "*". **Most Voted**

B. Use AWS Control Tower detective controls to block deployment of EC2 instances with public IP addresses and inline policies with elevated access or "*".

C. Use AWS Config to create rules for EC2 and IAM compliance. Configure the rules to run an AWS Systems Manager Session Manager automation to delete a resource when it is not compliant.

D. Use a service control policy (SCP) to block actions for the EC2 instances and IAM resources if the actions lead to noncompliance.

Correct Answer: *A*

Community vote distribution

A (70%) D (30%)

Comments

jaswantn Highly Voted 1 year, 3 months ago

Selected Answer: D

Option D... This is preventive control of Control Tower where we use SCP to disallow actions that lead to policy violation.

upvoted 9 times

[Removed] Highly Voted 1 year, 3 months ago

Selected Answer: A

proactive controls pls see links for both * in inline policy: https://docs.aws.amazon.com/controltower/latest/userguide/iam-rules.html#ct-iam-pr-1-description
and for ec2 public IP: https://docs.aws.amazon.com/controltower/latest/userguide/ec2-rules.html#ct-ec2-pr-9-description

upvoted 9 times

FlyingHawk Most Recent 4 months, 2 weeks ago

Selected Answer: A

Proactive controls check whether resources are compliant with your company policies and objectives, before the resources are provisioned in your accounts. If the resources are out of compliance, they are not provisioned. Proactive controls monitor resources that would be deployed in your accounts by means of AWS CloudFormation templates.

For those who are familiar with AWS: In AWS Control Tower preventive controls are implemented with Service Control Policies (SCPs).
Detective controls are implemented with AWS Config rules.
Proactive controls are implemented with AWS CloudFormation hooks.
upvoted 3 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: A

A - Proactive controls are a feature of AWS Control Tower that prevent noncompliant resources from being deployed by validating configurations before deployment.
B - Detective controls can't block deployment.
C - If you must monitor to get it right, then something's already wrong before you notice that.
D - SCPs can indeed block specific API calls for creating IAM resources with "*" or EC2 instances with public IPs, but can't make the most of AWS CloudFormation stacks. Plus, SCPs apply at the account level and might inadvertently restrict legitimate use cases. Not ideal enough is what I'm saying.
upvoted 1 times

LeonSauveterre 5 months, 2 weeks ago

And btw, since "The company has AWS Control Tower enabled", option A induces much less overhead than option D.
upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: A

Prevent AWS CloudFormation from deploying IAM resources and EC2 instances based on specific use cases = Control Tower Proactive controls.

"Proactive controls are security controls that are designed to prevent the creation of noncompliant resources.

For example (...), through AWS CloudFormation, the proactive control can prevent the creation of update of any S3 bucket that has public access enabled."

https://docs.aws.amazon.com/prescriptive-guidance/latest/aws-security-controls/proactive-controls.html
upvoted 3 times

88f8032 1 year, 1 month ago

Selected Answer: A

this is A
upvoted 2 times

Sergiuss95 1 year, 1 month ago

Selected Answer: D

Is D, the best way to prevent this actions, is deploying SCPs
upvoted 1 times

BBR01 1 year, 1 month ago

Selected Answer: D

It is D. You want to prevent the events from happening.
Proactive controls check whether resources are compliant with your company policies and objectives, before the resources are provisioned in your accounts.
Detective controls detect specific events when they occur and log the action in CloudTrail.
Preventive controls prevent actions from occurring.
Preventive controls are implemented with SCPs. Detective controls are implemented with AWS Config rules. Proactive controls are implemented with AWS CloudFormation hooks.
https://docs.aws.amazon.com/controltower/latest/userguide/how-control-tower-works.html#how-controls-work
upvoted 1 times

TwinSpark 1 year ago

as stated by you A is correct, Proactive controls are implemented as cloudformation hooks and the resource will not be deployed if not compliants. It is exactly what is asked in question.
Using an SCP it is actually a valid solution, but it need to be associated to a specific resource that is not specified. If you associated to root account nobody can deploy a public ip ec2, not only cloudformation

upvoted 2 times

osmk 1 year, 3 months ago

Selected Answer: A

Proactive controls are implemented using AWS CloudFormation hooks within AWS Control Tower. They operate before resources are deployed to determine compliance with activated controls. SCPs are part of AWS Organizations and are used to manage permissions. vs Define specific purposes for implementing controls.https://docs.aws.amazon.com/controltower/latest/userguide/proactive-controls.html

upvoted 6 times

osmk 1 year, 3 months ago

SCPs focus on managing permissions at the OU level, while proactive controls in AWS Control Tower help prevent non-compliance during resource provisioning.

upvoted 3 times

NayeraB 1 year, 3 months ago

Selected Answer: A

A would provide a proactive solution, also I'm not sure if SCP are made for granular details like creation of EC2 instances with public IP addresses or IAM resources with certain inline policies.

upvoted 2 times

Andy_09 1 year, 4 months ago

Option D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #790 *Topic 1*

A company's web application that is hosted in the AWS Cloud recently increased in popularity. The web application currently exists on a single Amazon EC2 instance in a single public subnet. The web application has not been able to meet the demand of the increased web traffic.

The company needs a solution that will provide high availability and scalability to meet the increased user demand without rewriting the web application.

Which combination of steps will meet these requirements? (Choose two.)

A. Replace the EC2 instance with a larger compute optimized instance.

B. Configure Amazon EC2 Auto Scaling with multiple Availability Zones in private subnets. **Most Voted**

C. Configure a NAT gateway in a public subnet to handle web requests.

D. Replace the EC2 instance with a larger memory optimized instance.

E. Configure an Application Load Balancer in a public subnet to distribute web traffic. **Most Voted**

Correct Answer: *BE*

Community vote distribution

BE (100%)

Comments

kempes Highly Voted 1 year, 4 months ago

Selected Answer: BE

Option BE

upvoted 6 times

Scheldon Most Recent 11 months, 3 weeks ago

Selected Answer: BE

AnswerBE,

AutoScaling to increase amount of servers per need,
Load Balancer to palance traffic equaly to all available servers

upvoted 2 times

upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: BE

Only BE makes sense, even though it might require modification of the application

upvoted 4 times

gsgdga 1 year, 2 months ago

Why isn't C the answer?

upvoted 1 times

802c4ff 1 year, 1 month ago

nat gateway is for accessing internet-facing from private subnets, not the other way around

upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: BE

be are correct

upvoted 2 times

chefKC 1 year, 4 months ago

Option B & E

upvoted 3 times

Andy_09 1 year, 4 months ago

Option BE

upvoted 4 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #791 *Topic 1*

A company has AWS Lambda functions that use environment variables. The company does not want its developers to see environment variables in plaintext.

Which solution will meet these requirements?

A. Deploy code to Amazon EC2 instances instead of using Lambda functions.

B. Configure SSL encryption on the Lambda functions to use AWS CloudHSM to store and encrypt the environment variables.

C. Create a certificate in AWS Certificate Manager (ACM). Configure the Lambda functions to use the certificate to encrypt the environment variables.

D. Create an AWS Key Management Service (AWS KMS) key. Enable encryption helpers on the Lambda functions to use the KMS key to store and encrypt the environment variables. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

osmk Highly Voted 1 year, 3 months ago

Selected Answer: D

https://docs.aws.amazon.com/lambda/latest/dg/configuration-envvars.html#configuration-envvars-encryption

upvoted 7 times

Andy_09 Highly Voted 1 year, 4 months ago

Option D

upvoted 7 times

JA2018 Most Recent 6 months, 1 week ago

Selected Answer: D

Reasons for choosing D:

AWS KMS is the primary service for managing encryption keys in AWS. This makes it the most appropriate choice for encrypting sensitive environment variables within Lambda functions.

Encryption helpers in Lambda allow you to easily integrate KMS key usage within your Lambda code, ensuring that environment variables are stored and accessed in an encrypted format.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

"To configure encryption for your environment variables

Enable console encryption helpers to use client-side encryption to protect your data in transit.

Under Encryption in transit, choose Enable helpers for encryption in transit.

For each environment variable that you want to enable console encryption helpers for, choose Encrypt next to the environment variable.

Under AWS KMS key to encrypt in transit, choose a customer managed key that you created at the beginning of this procedure."

upvoted 3 times

Rhydian25 11 months, 1 week ago

Selected Answer: D

I don't understand why we should use a complex way of encrypting variables instead of using Parameter Store... but in this case the best option is D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #792 *Topic 1*

An analytics company uses Amazon VPC to run its multi-tier services. The company wants to use RESTful APIs to offer a web analytics service to millions of users. Users must be verified by using an authentication service to access the APIs.

Which solution will meet these requirements with the MOST operational efficiency?

A. Configure an Amazon Cognito user pool for user authentication. Implement Amazon API Gateway REST APIs with a Cognito authorizer. **Most Voted**

B. Configure an Amazon Cognito identity pool for user authentication. Implement Amazon API Gateway HTTP APIs with a Cognito authorizer.

C. Configure an AWS Lambda function to handle user authentication. Implement Amazon API Gateway REST APIs with a Lambda authorizer.

D. Configure an IAM user to handle user authentication. Implement Amazon API Gateway HTTP APIs with an IAM authorizer.

Correct Answer: *A*

Community vote distribution

A (89%) B (11%)

Comments

stephensimudemy Highly Voted 9 months, 3 weeks ago

Selected Answer: A

User pools is for Authentication and user management

upvoted 7 times

sandordini Highly Voted 7 months, 2 weeks ago

Selected Answer: A

User pools are for authentication. Your app users can sign in through the user pool, Identity pools are for authorization, give them access to other AWS services.

upvoted 6 times

zdi561 Most Recent 3 months, 3 weeks ago

Selected Answer: A

When comparing operation efficiency, a Cognito authorizer in AWS API Gateway is generally considered more performant than a Lambda authorizer for basic authentication checks, as Cognito is a managed service optimized for user authentication, while a Lambda authorizer introduces an additional layer of processing with potential cold start overhead, making it better suited for complex, custom authorization logic. But Lambda is more flexible

upvoted 1 times

[Removed] 9 months, 1 week ago

Selected Answer: A

user pool vs identity pool: https://repost.aws/knowledge-center/cognito-user-pools-identity-pools

upvoted 3 times

NayeraB 9 months, 3 weeks ago

Selected Answer: B

B offers more operational efficiency imo

upvoted 2 times

chefKC 10 months ago

Answer is A

upvoted 2 times

Andy_09 10 months ago

Option A

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #793 *Topic 1*

A company has a mobile app for customers. The app's data is sensitive and must be encrypted at rest. The company uses AWS Key Management Service (AWS KMS).

The company needs a solution that prevents the accidental deletion of KMS keys. The solution must use Amazon Simple Notification Service (Amazon SNS) to send an email notification to administrators when a user attempts to delete a KMS key.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Amazon EventBridge rule that reacts when a user tries to delete a KMS key. Configure an AWS Config rule that cancels any deletion of a KMS key. Add the AWS Config rule as a target of the EventBridge rule. Create an SNS topic that notifies the administrators.

B. Create an AWS Lambda function that has custom logic to prevent KMS key deletion. Create an Amazon CloudWatch alarm that is activated when a user tries to delete a KMS key. Create an Amazon EventBridge rule that invokes the Lambda function when the DeleteKey operation is performed. Create an SNS topic. Configure the EventBridge rule to publish an SNS message that notifies the administrators.

C. Create an Amazon EventBridge rule that reacts when the KMS DeleteKey operation is performed. Configure the rule to initiate an AWS Systems Manager Automation runbook. Configure the runbook to cancel the deletion of the KMS key. Create an SNS topic. Configure the EventBridge rule to publish an SNS message that notifies the administrators. Most Voted

D. Create an AWS CloudTrail trail. Configure the trail to deliver logs to a new Amazon CloudWatch log group. Create a CloudWatch alarm based on the metric filter for the CloudWatch log group. Configure the alarm to use Amazon SNS to notify the administrators when the KMS DeleteKey operation is performed.

Correct Answer: *C*

Community vote distribution

C (86%) D (14%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option C
https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/monitor-and-remediate-scheduled-deletion-of-aws-kms-keys.html

upvoted 14 times

hajra313 Highly Voted 1 year, 4 months ago

option c bcz Option C emerges as the clear winner due to its:

Direct event monitoring for the DeleteKey operation

Pre-built automation using Systems Manager Automation runbooks

Efficient notification via Amazon SNS

Minimal code development and operational overhead

Reduced risk of accidental deletion with faster response times

upvoted 10 times

FlyingHawk Most Recent 4 months, 2 weeks ago

Selected Answer: C

A- AWS config is for detection, not for prevention, it cannot cancel the deletion.
B is more effort to implement the custom logic to prevent KMS key deletion.
C is correct to use EventBridge rule to detect it then config the rule to use AWS System Manager Automation rubook to cancel the deletion and publish the event to SNS.
D - does not have action for prevent the deletion of key.

upvoted 2 times

JA2018 6 months, 1 week ago

Selected Answer: C

Option C provides the most streamlined solution with minimal additional components needed (LEAST operational overhead). It directly leverages EventBridge to monitor for KMS key deletion attempts, triggers a Systems Manager Automation runbook to automatically prevent the deletion, and sends notifications through SNS without requiring additional Lambda functions or complex CloudWatch metric filtering.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: C

"Deletion of an AWS KMS key is scheduled.
The scheduled-deletion event is evaluated by an EventBridge rule.
The EventBridge rule engages the Amazon SNS topic.
The EventBridge rule initiates the Systems Manager automation and runbooks.
The runbooks cancel the deletion."

https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/monitor-and-remediate-scheduled-deletion-of-aws-kms-keys.html

upvoted 2 times

JunsK1e 11 months ago

Selected Answer: C

I agree with andy_09

upvoted 2 times

Dammy031 11 months, 2 weeks ago

Selected Answer: D

Cloud trail helps to keep all invoked API calls in the AWS account which can trail back to the delete call made by a user
CloudWatch triggers an alarm when deletion is attempted.
SNS sends a notification to the administration about the attempt made.

All these met the requirement of the question.

upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: C

My educated guess was C. Now, reading the comments, from Hajra313 and knben I feel confident as well :)
upvoted 3 times

camps 1 year, 2 months ago

It's D https://docs.aws.amazon.com/kms/latest/developerguide/deleting-keys-creating-cloudwatch-alarm.html

https://docs.aws.amazon.com/kms/latest/developerguide/deleting-keys-creating-cloudwatch-alarm.html#cloudwatch-alarm-prerequisites
upvoted 2 times

Salilgen 5 months, 1 week ago

This procedures create an alarm that notifies you when someone in your account tries to use a KMS key that is pending deletion. It does't notify you when a user attempts to delete a KMS key or even prevent deletion of the key
upvoted 1 times

1dd 1 year, 3 months ago

C as it " cancel the deletion of the KMS key"
upvoted 2 times

knben 1 year, 3 months ago

I would go with C

A -> Config is for compliance
B -> No lambda is required, too much complexity
C -> It achieves the goal, since KMS keys are not immediately deleted, which gives time to automation to cancel the action
D -> Cloudtrail is for auditing
upvoted 4 times

NayeraB 1 year, 3 months ago

Selected Answer: C

I agree with hajra313
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #794 *Topic 1*

A company wants to analyze and generate reports to track the usage of its mobile app. The app is popular and has a global user base. The company uses a custom report building program to analyze application usage.

The program generates multiple reports during the last week of each month. The program takes less than 10 minutes to produce each report. The company rarely uses the program to generate reports outside of the last week of each month The company wants to generate reports in the least amount of time when the reports are requested.

Which solution will meet these requirements MOST cost-effectively?

A. Run the program by using Amazon EC2 On-Demand Instances. Create an Amazon EventBridge rule to start the EC2 instances when reports are requested. Run the EC2 instances continuously during the last week of each month.

B. Run the program in AWS Lambda. Create an Amazon EventBridge rule to run a Lambda function when reports are requested. **Most Voted**

C. Run the program in Amazon Elastic Container Service (Amazon ECS). Schedule Amazon ECS to run the program when reports are requested.

D. Run the program by using Amazon EC2 Spot Instances. Create an Amazon EventBndge rule to start the EC2 instances when reports are requested. Run the EC2 instances continuously during the last week of each month.

Correct Answer: *B*

Community vote distribution

B (82%) C (18%)

Comments

Andy_09 Highly Voted 1 year, 4 months ago

Option B

upvoted 6 times

LuongTo 7 months, 2 weeks ago

can lambda host the "custom report building program" ?

upvoted 2 times

ogerber 1 year, 3 months ago

not sure because it says that the program producues several reports , and each takes less than 10 min. i am voting for option A
upvoted 2 times

1dd 1 year, 3 months ago

Lambda takes duration--> 15 minutes
upvoted 1 times

FZA24 1 year, 3 months ago

each lambda triggering produces a report in less than 10 mins.
upvoted 4 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: B

A - The task only runs for 10 minutes per report. Why would you need an EC2 instance?
B - No infrastructure to manage. Love it.
C - Works, but cost more than B and is more complex to set up than B.
D - More cost-effective than A, but spot instances can be interrupted, probably causing delays in report generation.
upvoted 1 times

agbor_tambe 8 months, 1 week ago

Selected Answer: B

B is correct
upvoted 2 times

EdricHoang 11 months, 2 weeks ago

Selected Answer: B

"The program takes less than 10 minutes to produce each report. The company rarely uses the program to generate reports outside of the last week of each month The company wants to generate reports in the least amount of time when the reports are requested."
I go for B because of this
upvoted 2 times

Scheldon 12 months ago

Selected Answer: C

AnswerC

Option A and D are saying abount Running the EC2 instances countinously during the last week of each month which is not necessery from my understanding and will be not so cheap.

Option B - 10min per report and we have couple of reports. So it looks like program is running for at least 20 min so in theory Lambda is not useful here

Option C - ECS is allowing us to run Fargate which will allow to run program for more then 15 min, hence all reports which program is preparing should be created. I'm not sure but I think ECS API is allowing to run task on demand/request.
upvoted 2 times

MatAlves 8 months, 3 weeks ago

It's "15 minutes per execution", which is enough to produce a report and be ready for another invocation.
upvoted 2 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: B

B is corre ct
upvoted 2 times

NayeraB 1 year, 3 months ago

Selected Answer: B

B..maybe?

upvoted 2 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #795 *Topic 1*

A company is designing a tightly coupled high performance computing (HPC) environment in the AWS Cloud. The company needs to include features that will optimize the HPC environment for networking and storage.

Which combination of solutions will meet these requirements? (Choose two.)

A. Create an accelerator in AWS Global Accelerator. Configure custom routing for the accelerator.

B. Create an Amazon FSx for Lustre file system. Configure the file system with scratch storage. Most Voted

C. Create an Amazon CloudFront distribution. Configure the viewer protocol policy to be HTTP and HTTPS.

D. Launch Amazon EC2 instances. Attach an Elastic Fabric Adapter (EFA) to the instances. Most Voted

E. Create an AWS Elastic Beanstalk deployment to manage the environment.

Correct Answer: *BD*

Community vote distribution

BD (100%)

Comments

Andy_09 Highly Voted 10 months ago

Options BD

upvoted 7 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: BD

A - Global Accelerator is not designed for tightly coupled HPC environments. It focuses on improving performance for geographically distributed users, not for low-latency intra-cluster communication.
B - FSx for Lustre is specifically designed for HPC workloads and integrates well with Amazon S3 for high-throughput and low-latency access to data.
C - CloudFront does not address the low-latency, high-bandwidth networking or storage requirements of tightly coupled HPC environments.
D - EFA is a network interface that provides low-latency, high-throughput connectivity for tightly coupled HPC workloads, such as those using MPI (Message Passing Interface).
E - Elastic Beanstalk is used for deploying and managing web applications and services, which has nothing to do with HPC.

upvoted 2 times

seetpt 9 months ago

Selected Answer: BD

BD looks right

upvoted 3 times

asdfcdsxdfc 9 months ago

Selected Answer: BD

Elastic Fabric Adapter (EFA)
• Improved ENA for HPC, only works for Linux
• Great for inter-node communications, tightly coupled workloads
• Leverages Message Passing Interface (MPI) standard
• Bypasses the underlying Linux OS to provide low-latency, reliable transport

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #796 *Topic 1*

A company needs a solution to prevent photos with unwanted content from being uploaded to the company's web application. The solution must not involve training a machine learning (ML) model.

Which solution will meet these requirements?

A. Create and deploy a model by using Amazon SageMaker Autopilot. Create a real-time endpoint that the web application invokes when new photos are uploaded.

B. Create an AWS Lambda function that uses Amazon Rekognition to detect unwanted content. Create a Lambda function URL that the web application invokes when new photos are uploaded. **Most Voted**

C. Create an Amazon CloudFront function that uses Amazon Comprehend to detect unwanted content. Associate the function with the web application.

D. Create an AWS Lambda function that uses Amazon Rekognition Video to detect unwanted content. Create a Lambda function URL that the web application invokes when new photos are uploaded.

Correct Answer: *B*

Community vote distribution

B (93%) A (7%)

Comments

Andy_09 Highly Voted 10 months ago

Option B

upvoted 8 times

HarryLopez Highly Voted 9 months ago

Selected Answer: B

Rekognition: for image and video analysis
Comprehend: natural language processing model for uncovering insights and connections in text
Sagemaker Autopilot: feature set that simplifies and accelerates and automates the various stages of the machine learning workflow

upvoted 6 times

zdi561 Most Recent 3 months, 3 weeks ago

Selected Answer: A

Amazon Rekognition is a visual analysis service that includes both Amazon Rekognition Image and Amazon Rekognition Video. Amazon Rekognition Image analyzes images, while Amazon Rekognition Video analyzes videos.
upvoted 1 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: B

You've made it this far, and you're all gonna make it.

For your quick reference:
Vision and Image Processing
1. Amazon Rekognition: Detects objects, scenes, text, faces, and activities in images/videos.
2. Amazon Textract: Extracts text, tables, and forms from scanned documents.

Speech and Language
1. Amazon Polly: Converts text to lifelike speech.
2. Amazon Transcribe: Converts speech to text for audio files.
3. Amazon Translate: Yeah, translate.
4. Amazon Comprehend: Extracts insights like sentiment, entities, and key phrases from text.

Machine Learning Development
1. Amazon SageMaker: Comprehensive ML platform for building, training, and deploying models.
2. Amazon SageMaker Autopilot: Automatically builds, trains, and tunes ML models, basically simplified model creation for non-experts.
3. Amazon SageMaker JumpStart: Provides pre-built ML models and solutions for quick deployment.
upvoted 5 times

LeonSauveterre 5 months, 2 weeks ago

Chatbots and Conversations
1. Amazon Lex: Builds conversational AI (chatbots) using speech and text.
2. Amazon Kendra: Intelligent search service for documents and FAQs.

Personalization and Forecasting
1. Amazon Personalize: Provides personalized recommendations using ML.
2. Amazon Forecast: Time-series forecasting for demand, sales, and resource planning.

Data Insights
1. Amazon Fraud Detector: Detects fraudulent activity using ML models.
2. Amazon Lookout for Metrics: Detects anomalies in metrics automatically.
3. Amazon Lookout for Vision: Detects manufacturing defects in images.

Robotics and Edge AI
1. AWS DeepRacer: Provides a platform for learning reinforcement learning using autonomous racing.
2. AWS Panorama: Adds computer vision to on-premises cameras.
upvoted 4 times

NayeraB 9 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #797 *Topic 1*

A company uses AWS to run its ecommerce platform. The platform is critical to the company's operations and has a high volume of traffic and transactions. The company configures a multi-factor authentication (MFA) device to secure its AWS account root user credentials. The company wants to ensure that it will not lose access to the root user account if the MFA device is lost.

Which solution will meet these requirements?

A. Set up a backup administrator account that the company can use to log in if the company loses the MFA device.

B. Add multiple MFA devices for the root user account to handle the disaster scenario. Most Voted

C. Create a new administrator account when the company cannot access the root account.

D. Attach the administrator policy to another IAM user when the company cannot access the root account.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

hajra313 Highly Voted 1 year, 4 months ago

B. Add multiple MFA devices for the root user account to handle the disaster scenario.

By adding multiple MFA devices for the root user account, the company ensures that it can still access the account even if one MFA device is lost. This approach provides a backup for authentication, addressing the concern of losing access to the root user account if the MFA device is lost.

upvoted 11 times

Tatai2015 1 year ago

https://docs.aws.amazon.com/IAM/latest/UserGuide/root-user-best-practices.html

upvoted 2 times

ad2dj28 Most Recent 3 months ago

Selected Answer: B

Your AWS account root user and IAM users can register up to eight MFA devices of any type. Registering multiple MFA devices can provide flexibility and help you reduce the risk of access interruption if a device is lost or broken.

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_credentials_mfa.html

upvoted 1 times

Scheldon 12 months ago

Selected Answer: B

AnswerB

Because a root user can perform privileged actions, it's crucial to add MFA for the root user as a second authentication factor in addition to the email address and password as sign-in credentials. We strongly recommend enabling multiple MFA for your root user credentials to provide additional flexibility and resiliency in your security strategy. You can register up to eight MFA devices of any combination of the currently supported MFA types with your AWS account root user.
https://docs.aws.amazon.com/IAM/latest/UserGuide/root-user-best-practices.html#ru-bp-mfa

upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: B

b looks correct

upvoted 2 times

NayeraB 1 year, 3 months ago

Selected Answer: B

I'd go for B

upvoted 3 times

Andy_09 1 year, 4 months ago

Option B

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #798 *Topic 1*

A social media company is creating a rewards program website for its users. The company gives users points when users create and upload videos to the website. Users redeem their points for gifts or discounts from the company's affiliated partners. A unique ID identifies users. The partners refer to this ID to verify user eligibility for rewards.

The partners want to receive notification of user IDs through an HTTP endpoint when the company gives users points. Hundreds of vendors are interested in becoming affiliated partners every day. The company wants to design an architecture that gives the website the ability to add partners rapidly in a scalable way.

Which solution will meet these requirements with the LEAST implementation effort?

A. Create an Amazon Timestream database to keep a list of affiliated partners. Implement an AWS Lambda function to read the list. Configure the Lambda function to send user IDs to each partner when the company gives users points.

B. Create an Amazon Simple Notification Service (Amazon SNS) topic. Choose an endpoint protocol. Subscribe the partners to the topic. Publish user IDs to the topic when the company gives users points. Most Voted

C. Create an AWS Step Functions state machine. Create a task for every affiliated partner. Invoke the state machine with user IDs as input when the company gives users points.

D. Create a data stream in Amazon Kinesis Data Streams. Implement producer and consumer applications. Store a list of affiliated partners in the data stream. Send user IDs when the company gives users points.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

kempes Highly Voted 1 year, 4 months ago

Selected Answer: B

SNS is designed for precisely this kind of use case. It allows you to publish messages to a topic, which can then be delivered to multiple subscribers. Partners can subscribe to the SNS topic using an HTTP endpoint as the protocol, which meets the requirement to notify partners via an HTTP endpoint. This approach is highly scalable and requires the least implementation effort because it leverages managed services without the need for custom logic to manage subscriptions or deliver notifications.

upvoted 13 times

hajra313 Highly Voted 1 year, 4 months ago

Option A involves creating an Amazon Timestream database to store affiliated partners and implementing an AWS Lambda function to read the list and send user IDs to each partner. While this approach can work, it involves more implementation effort than the Amazon SNS solution. It requires setting up and managing a database, as well as configuring the Lambda function to send notifications to partners. The Amazon SNS solution provides a simpler and more scalable approach for rapidly adding partners and notifying them when users receive points. so answer is B

upvoted 9 times

Scheldon Most Recent 12 months ago

Selected Answer: B

AnswerB

Sending Notification to multiple subscribers = SNS

Amazon Simple Notification Service (Amazon SNS) is a managed service that provides message delivery from publishers to subscribers (also known as producers and consumers). Publishers communicate asynchronously with subscribers by sending messages to a topic, which is a logical access point and communication channel. Clients can subscribe to the SNS topic and receive published messages using a supported endpoint type, such as Amazon Data Firehose, Amazon SQS, AWS Lambda, HTTP, email, mobile push notifications, and mobile text messages (SMS).
https://docs.aws.amazon.com/sns/latest/dg/welcome.html

upvoted 1 times

NayeraB 1 year, 3 months ago

Selected Answer: B

This is a perfect SNS use case

upvoted 2 times

jjcode 1 year, 4 months ago

The answer is B, create an SNS topic one subscriptions you can make is HTTP, This completely addresses the question objective.

upvoted 2 times

Andy_09 1 year, 4 months ago

Option A

upvoted 1 times

FlyingHawk 4 months, 1 week ago

Timestream is optimized for time-series data, not for managing partner subscriptions.

upvoted 1 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #799 *Topic 1*

A company needs to extract the names of ingredients from recipe records that are stored as text files in an Amazon S3 bucket. A web application will use the ingredient names to query an Amazon DynamoDB table and determine a nutrition score.

The application can handle non-food records and errors. The company does not have any employees who have machine learning knowledge to develop this solution.

Which solution will meet these requirements MOST cost-effectively?

A. Use S3 Event Notifications to invoke an AWS Lambda function when PutObject requests occur. Program the Lambda function to analyze the object and extract the ingredient names by using Amazon Comprehend. Store the Amazon Comprehend output in the DynamoDB table. **Most Voted**

B. Use an Amazon EventBridge rule to invoke an AWS Lambda function when PutObject requests occur. Program the Lambda function to analyze the object by using Amazon Forecast to extract the ingredient names. Store the Forecast output in the DynamoDB table.

C. Use S3 Event Notifications to invoke an AWS Lambda function when PutObject requests occur. Use Amazon Polly to create audio recordings of the recipe records. Save the audio files in the S3 bucket. Use Amazon Simple Notification Service (Amazon SNS) to send a URL as a message to employees. Instruct the employees to listen to the audio files and calculate the nutrition score. Store the ingredient names in the DynamoDB table.

D. Use an Amazon EventBridge rule to invoke an AWS Lambda function when a PutObject request occurs. Program the Lambda function to analyze the object and extract the ingredient names by using Amazon SageMaker. Store the inference output from the SageMaker endpoint in the DynamoDB table.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

seetpt Highly Voted 1 year, 3 months ago

Selected Answer: A

A correct

upvoted 5 times

upvoted 5 times

Besisco Most Recent 3 months, 4 weeks ago

Selected Answer: A

I'm fascinated by how self-descriptive the names of AWS ML tools are.

upvoted 1 times

Danilus 6 months, 3 weeks ago

Selected Answer: A

Key: Most cost-effective solution
Key: Extract the names of ingredients from recipe records
Key: The company does not have any employees with machine learning knowledge

Not B because Forecast is used for making predictions in time series, not for extracting text or ingredients.
Not C because it would be too expensive and involve too much operational overhead.
Not D because SageMaker is designed for creating custom machine learning models, and the company lacks employees with machine learning knowledge.

Conclusion: The answer is A because Comprehend is used for natural language processing (NLP), which in this case can extract ingredient names from the recipes effectively.

so the answer is A because comprehend is used for nlp (natular languages processing) which is in this case use for to extract the names of the recipes

upvoted 3 times

Scheldon 12 months ago

Selected Answer: A

AnswerA

Amazon Comprehend uses natural language processing (NLP) to extract insights about the content of documents. It develops insights by recognizing the entities, key phrases, language, sentiments, and other common elements in a document. Use Amazon Comprehend to create new products based on understanding the structure of documents. For example, using Amazon Comprehend you can search social networking feeds for mentions of products or scan an entire document repository for key phrases.
https://docs.aws.amazon.com/comprehend/latest/dg/what-is.html

upvoted 3 times

fae3297 1 year ago

Selected Answer: A

A seems right

B. Forecast is time-series
C. Using Polly to create audio recordings just to make your employees listen to them seems inefficient to say the least
D. Question asks for no ML. SageMaker = ML

upvoted 3 times

seetpt 1 year, 3 months ago

Selected Answer: A

A is correct

upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

shouldn't it be A?

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #800 *Topic 1*

A company needs to create an AWS Lambda function that will run in a VPC in the company's primary AWS account. The Lambda function needs to access files that the company stores in an Amazon Elastic File System (Amazon EFS) file system. The EFS file system is located in a secondary AWS account. As the company adds files to the file system, the solution must scale to meet the demand.

Which solution will meet these requirements MOST cost-effectively?

A. Create a new EFS file system in the primary account. Use AWS DataSync to copy the contents of the original EFS file system to the new EFS file system.

B. Create a VPC peering connection between the VPCs that are in the primary account and the secondary account. **Most Voted**

C. Create a second Lambda function in the secondary account that has a mount that is configured for the file system. Use the primary account's Lambda function to invoke the secondary account's Lambda function.

D. Move the contents of the file system to a Lambda layer. Configure the Lambda layer's permissions to allow the company's secondary account to use the Lambda layer.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

lenotc Highly Voted 1 year, 3 months ago

Selected Answer: B

B -> VPC peering allows the Lambda access secondary account securely and efficiently
A -> redundancy
C -> additional complexity
D -> sharing code libraries

upvoted 9 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: B

A - Duplicating the EFS file system leads to unnecessary storage costs and additional complexity.

A - Duplicating the EFS file system leads to unnecessary storage costs and additional complexity.
B - VPC peering connection allows seamless access without additional manual scaling.
C - Well... I don't know, but you could've directly enabled access to the EFS file system. Why make it even more complex?
D - Lambda layers are intended for shared libraries or dependencies. You would be crazy to store large amounts of data there.

upvoted 1 times

Scheldon 12 months ago

Selected Answer: B

AnswerB

You can configure a function to mount an Amazon EFS file system in another AWS account. Before you mount the file system, you must ensure the following:
VPC peering must be configured, and appropriate routes must be added to the route tables in each VPC.
.
.
.
https://docs.aws.amazon.com/lambda/latest/dg/configuration-filesystem-cross-account.html

upvoted 1 times

Nm55569 1 year ago

Selected Answer: B

https://docs.aws.amazon.com/lambda/latest/dg/configuration-filesystem.html#configuration-filesystem-cross-account

upvoted 1 times

osmk 1 year, 3 months ago

Selected Answer: B

https://docs.aws.amazon.com/efs/latest/ug/efs-different-vpc.html

upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

Shouldn't it be B?

upvoted 1 times

1dd 1 year, 3 months ago

I thinks AWS DataSync less costly

upvoted 1 times

[Removed] 1 year, 2 months ago

setting up a peering connection is free. same for data transfer in the same AZ. data sync at the end of the day cost $$$ to move data.

upvoted 3 times

Scheldon 12 months ago

When you will SyncData you need to have secondary storage for which you need to pay so it is not cheap solution.

upvoted 1 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #801 *Topic 1*

A financial company needs to handle highly sensitive data. The company will store the data in an Amazon S3 bucket. The company needs to ensure that the data is encrypted in transit and at rest. The company must manage the encryption keys outside the AWS Cloud.

Which solution will meet these requirements?

A. Encrypt the data in the S3 bucket with server-side encryption (SSE) that uses an AWS Key Management Service (AWS KMS) customer managed key.

B. Encrypt the data in the S3 bucket with server-side encryption (SSE) that uses an AWS Key Management Service (AWS KMS) AWS managed key.

C. Encrypt the data in the S3 bucket with the default server-side encryption (SSE).

D. Encrypt the data at the company's data center before storing the data in the S3 bucket. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

srcntpc 3 weeks, 1 day ago

Selected Answer: D

D for sure

upvoted 1 times

Johnoppong101 9 months, 3 weeks ago

You get to do it, keep moving...

upvoted 4 times

Scheldon 12 months ago

Selected Answer: D

AnswerD

Hence we need to encrypt data not only during the rest but during the transfer as well, we need execute client-side encyprion. SSE will only secure data during rest hence we can eliminate A,B and C.
https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingClientSideEncryption.html

upvoted 1 times

chasingsummer 1 year, 2 months ago

Selected Answer: D

Given the requirement to manage encryption keys outside the AWS Cloud, option D is the most suitable solution, despite not directly utilizing AWS's native encryption services like SSE with AWS KMS. Instead, it leverages external encryption mechanisms controlled by the company.

upvoted 4 times

rondelldell 1 year, 2 months ago

A Key is safe but came from the customer

upvoted 2 times

Mikado211 1 year, 2 months ago

Selected Answer: D

A, B and C need to have the key stored in AWS cloud.
D is correct.

upvoted 3 times

osmk 1 year, 3 months ago

Selected Answer: D

Client-side encryption – You encrypt your data client-side and upload the encrypted data to Amazon S3. In this case, you manage the encryption process, encryption keys, and related tools.https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingClientSideEncryption.html

upvoted 3 times

giovanna_mag 1 year, 3 months ago

Selected Answer: D

For me it's D, it's the only one that provides encryption also in transit

upvoted 2 times

asdfcdsxdfc 1 year, 3 months ago

A looks correct

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #802 *Topic 1*

A company wants to run its payment application on AWS. The application receives payment notifications from mobile devices. Payment notifications require a basic validation before they are sent for further processing.

The backend processing application is long running and requires compute and memory to be adjusted. The company does not want to manage the infrastructure.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Amazon Simple Queue Service (Amazon SQS) queue. Integrate the queue with an Amazon EventBridge rule to receive payment notifications from mobile devices. Configure the rule to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon Elastic Kubernetes Service (Amazon EKS) Anywhere. Create a standalone cluster.

B. Create an Amazon API Gateway API. Integrate the API with an AWS Step Functions state machine to receive payment notifications from mobile devices. Invoke the state machine to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon Elastic Kubernetes Service (Amazon EKS). Configure an EKS cluster with self-managed nodes.

C. Create an Amazon Simple Queue Service (Amazon SQS) queue. Integrate the queue with an Amazon EventBridge rule to receive payment notifications from mobile devices. Configure the rule to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon EC2 Spot Instances. Configure a Spot Fleet with a default allocation strategy.

D. Create an Amazon API Gateway API. Integrate the API with AWS Lambda to receive payment notifications from mobile devices. Invoke a Lambda function to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon Elastic Container Service (Amazon ECS). Configure Amazon ECS with an AWS Fargate launch type. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

Mikado211 Highly Voted 1 year, 2 months ago

Selected Answer: D

We want to have least overhead and no infrastructure (aka no server).
So no infrastructure == not C
least overhead == ECS better than EKS == not B and not A

Fargate is serverless so D is still valid.

So the answer is D.

upvoted 5 times

Cpso Most Recent 6 months, 1 week ago

Selected Answer: D

D= less operation. But should have SQS to buffer a job.

upvoted 2 times

Scheldon 12 months ago

Selected Answer: D

AnswerD

I would go with Lambda function to do basic validation it should not take more then 15min hence Lambda is perfect for that job. Then we have information that backend processing application is long running and we need to make compute and memory adjustment, and everythink need to be automatic as company do not want to manage infrastucture. In that situation Fargate with ECS will be ideal as it can run background application for every paymant separetly we need only adjust amunt of resources in use. More payments more application running, mmore resources in use and oposite,

upvoted 1 times

sandordini 1 year, 1 month ago

Selected Answer: D

Lot of grip in this question, but to keep it short:
No infra: B, C Out.
EKS Anywhere: Onprem + AWS: Not needed.
ECS Fargate: Serverless, Least Ops Overhead, SQS fine for the queue, Lambda good for basic validation.

upvoted 3 times

seetpt 1 year, 3 months ago

Selected Answer: D

D is correct

upvoted 4 times

asdfcdsxdfc 1 year, 3 months ago

shouldn't it be D?

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #803 *Topic 1*

A solutions architect is designing a user authentication solution for a company. The solution must invoke two-factor authentication for users that log in from inconsistent geographical locations, IP addresses, or devices. The solution must also be able to scale up to accommodate millions of users.

Which solution will meet these requirements?

A. Configure Amazon Cognito user pools for user authentication. Enable the risk-based adaptive authentication feature with multifactor authentication (MFA). **Most Voted**

B. Configure Amazon Cognito identity pools for user authentication. Enable multi-factor authentication (MFA).

C. Configure AWS Identity and Access Management (IAM) users for user authentication. Attach an IAM policy that allows the AllowManageOwnUserMFA action.

D. Configure AWS IAM Identity Center (AWS Single Sign-On) authentication for user authentication. Configure the permission sets to require multi-factor authentication (MFA).

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

mk168898 7 months, 3 weeks ago

Selected Answer: A

Picked A because cognito user pool => authentication

upvoted 1 times

Scheldon 12 months ago

Selected Answer: A

AnswerA

Adaptive authentication
Amazon Cognito can review location and device information from your users' sign-in requests and apply an automatic response to secure the user accounts in your user pool against suspicious activity.

When you activate advanced security, Amazon Cognito assigns a risk score to user activity. You can assign an automatic

When you activate advanced security, Amazon Cognito assigns a risk score to user activity. You can assign an automatic response to suspicious activity: you can Require MFA, Block sign-in, or just log the activity details and risk score. You can also automatically send email messages that notify your user of the suspicious activity so that they can reset their password or take other self-guided actions.
https://docs.aws.amazon.com/cognito/latest/developerguide/cognito-user-pool-settings-adaptive-authentication.html
upvoted 3 times

osmk 1 year, 3 months ago

Selected Answer: A

With adaptive authentication, you can configure your user pool to require second factor authentication in response to an increased risk level. To add adaptive authentication to your user pool, see Adding advanced security to a user pool.https://docs.aws.amazon.com/cognito/latest/developerguide/cognito-user-pool-settings-advanced-security.html
upvoted 3 times

lenotc 1 year, 3 months ago

Selected Answer: A

A is correct
B is wrong beacause is desing for temporary credentials
upvoted 2 times

giovanna_mag 1 year, 3 months ago

Selected Answer: A

I belive it's A
upvoted 1 times

Tatai2015 1 year ago

A
https://docs.aws.amazon.com/cognito/latest/developerguide/cognito-user-pool-settings-adaptive-authentication.html
upvoted 1 times

xBUGx 1 year, 3 months ago

accommodate millions of users and GEO, IP, etc. I think A
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #804 *Topic 1*

A company has an Amazon S3 data lake. The company needs a solution that transforms the data from the data lake and loads the data into a data warehouse every day. The data warehouse must have massively parallel processing (MPP) capabilities.

Data analysts then need to create and train machine learning (ML) models by using SQL commands on the data. The solution must use serverless AWS services wherever possible.

Which solution will meet these requirements?

A. Run a daily Amazon EMR job to transform the data and load the data into Amazon Redshift. Use Amazon Redshift ML to create and train the ML models.

B. Run a daily Amazon EMR job to transform the data and load the data into Amazon Aurora Serverless. Use Amazon Aurora ML to create and train the ML models.

C. Run a daily AWS Glue job to transform the data and load the data into Amazon Redshift Serverless. Use Amazon Redshift ML to create and train the ML models. **Most Voted**

D. Run a daily AWS Glue job to transform the data and load the data into Amazon Athena tables. Use Amazon Athena ML to create and train the ML models.

Correct Answer: *C*

Community vote distribution

C (90%) B (10%)

Comments

Mikado211 Highly Voted 1 year, 2 months ago

Selected Answer: C

Data warehouse ==> Redshift
Without additional informations both EMR and Glue Jobs can work.
Since the question asks to use serverless as much as possible, Redshift Serverless is a better solution.
C

upvoted 6 times

BatVanyo Highly Voted 1 year, 2 months ago

Selected Answer: C

Neither A, nor B explicitly say "EMR serverless" which is a new AWS offering, so I exclude these two.
MPP goes hand in hand with Redshift, so D is also incorrect.

This leaves C the only possible serverless option here.
upvoted 5 times

sk1974 Most Recent 3 months, 1 week ago

Selected Answer: B

Any reason it cannot be B ?
upvoted 2 times

mk168898 7 months, 3 weeks ago

Selected Answer: C

Data Warehouse => redshift
Use AWS Services whereever possible => Redshift serverless
upvoted 2 times

rondelldell 1 year, 2 months ago

A
Amazon EMR Serverless is a deployment option for Amazon EMR that provides a serverless runtime environment. This simplifies the operation of analytics applications that use the latest open-source frameworks, such as Apache Spark and Apache Hive. With EMR Serverless, you don't have to configure, optimize, secure, or operate clusters to run applications with these frameworks.

EMR Serverless helps you avoid over- or under-provisioning resources for your data processing jobs. EMR Serverless automatically determines the resources that the application needs, gets these resources to process your jobs, and releases the resources when the jobs finish. For use cases where applications need a response within seconds, such as interactive data analysis, you can pre-initialize the resources that the application needs when you create the application.
upvoted 2 times

1dd 1 year, 3 months ago

Selected Answer: C

Option C
upvoted 2 times

1dd 1 year, 3 months ago

EMR works with big data transfer
upvoted 1 times

1dd 1 year, 3 months ago

MPP --> use Redshift so eliminate B,D
As it required Serverless services --> Glue
upvoted 1 times

1dd 1 year, 3 months ago

A have no serverless
C is the answer
upvoted 1 times

seetpt 1 year, 3 months ago

Selected Answer: C

C is correct
upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

should be C
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #805 *Topic 1*

A company runs containers in a Kubernetes environment in the company's local data center. The company wants to use Amazon Elastic Kubernetes Service (Amazon EKS) and other AWS managed services. Data must remain locally in the company's data center and cannot be stored in any remote site or cloud to maintain compliance.

Which solution will meet these requirements?

A. Deploy AWS Local Zones in the company's data center.

B. Use an AWS Snowmobile in the company's data center.

C. Install an AWS Outposts rack in the company's data center. **Most Voted**

D. Install an AWS Snowball Edge Storage Optimized node in the data center.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

Mikado211 Highly Voted 1 year, 2 months ago

Selected Answer: C

Outpost is a service where AWS has physical servers in your datacenter.
C

upvoted 8 times

asdfcdsxdfc Highly Voted 1 year, 3 months ago

Selected Answer: C

C looks correct

upvoted 5 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: C

A - AWS Local Zones provide low-latency access to AWS services by placing AWS infrastructure closer to customers, but the data is still stored in an AWS-managed facility.
B - AWS Snowmobile is used for large-scale data migrations to the cloud, not for running workloads or storing data locally.
C - AWS Outposts allows you to run AWS services including Amazon EKS, in your own data center while keeping data locally.

C - AWS Outposts allows you to run AWS services, including Amazon EKS, in your own data center while keeping data locally.
D - AWS Snowball Edge is a portable edge computing and data storage device, but it's not suitable for running Kubernetes or AWS-managed services like Amazon EKS.

upvoted 2 times

mk168898 7 months, 3 weeks ago

Data must remain locally in the company's data center => AWS outpost

upvoted 2 times

JonJon03 11 months, 3 weeks ago

EKS on SnowBALL could be an option. EKS on SnowMOBILE isn't as it's used for data transfer mostly.

upvoted 2 times

Scheldon 12 months ago

Selected Answer: C

AnswerC

AWS Outpost = Bring AWS cloud to your DataCentre, which we need in described scenario

upvoted 4 times

seetpt 1 year, 3 months ago

Selected Answer: C

C is correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #806 *Topic 1*

A social media company has workloads that collect and process data. The workloads store the data in on-premises NFS storage. The data store cannot scale fast enough to meet the company's expanding business needs. The company wants to migrate the current data store to AWS.

Which solution will meet these requirements MOST cost-effectively?

A. Set up an AWS Storage Gateway Volume Gateway. Use an Amazon S3 Lifecycle policy to transition the data to the appropriate storage class.

B. Set up an AWS Storage Gateway Amazon S3 File Gateway. Use an Amazon S3 Lifecycle policy to transition the data to the appropriate storage class. **Most Voted**

C. Use the Amazon Elastic File System (Amazon EFS) Standard-Infrequent Access (Standard-IA) storage class. Activate the infrequent access lifecycle policy.

D. Use the Amazon Elastic File System (Amazon EFS) One Zone-Infrequent Access (One Zone-IA) storage class. Activate the infrequent access lifecycle policy.

Correct Answer: *B*

Community vote distribution

B (96%) D

Comments

alawada Highly Voted 1 year, 2 months ago

Selected Answer: B

This solution meets the requirements most cost-effectively because it enables the company to migrate its on-premises NFS data store to AWS without changing the existing applications or workflows. AWS Storage Gateway is a hybrid cloud storage service that provides seamless and secure integration between on-premises and AWS storage. Amazon S3 File Gateway is a type of AWS Storage Gateway that provides a file interface to Amazon S3, with local caching for low-latency access. By setting up an Amazon S3 File Gateway, the company can store and retrieve files as objects in Amazon S3 using standard file protocols such as NFS.

upvoted 9 times

0c21057 Most Recent 1 month, 3 weeks ago

Selected Answer: B

Why other options are incorrect:

A. Volume Gateway is designed for block storage and backup use cases, not for NFS file storage migration. It would be less cost-effective and more complex to implement.

C. Amazon EFS Standard-IA would be more expensive than S3 for this use case, as it's designed for file systems that need high availability across multiple AZs.

D. Amazon EFS One Zone-IA, while cheaper than Standard-IA, is still more expensive than S3 for this type of data storage requirement. It's also designed for file systems rather than object storage.

The S3 File Gateway solution is most cost-effective because:

S3 storage is generally cheaper than EFS
Lifecycle policies can automatically move data to cheaper storage tiers
It maintains NFS compatibility while providing scalability
It minimizes the need for application changes during migration
upvoted 3 times

Gorkha 4 months, 1 week ago

Selected Answer: B

as their is nfs = file gateway
upvoted 1 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: B

I'm a little confused by option C & D. How does infrequent access have anything to do with this?
upvoted 1 times

mk168898 7 months, 3 weeks ago

Selected Answer: B

S3 File Gateway => Best for NFS-like file storage workloads
upvoted 2 times

hharbiordun85 8 months ago

Answer: C
upvoted 1 times

Scheldon 12 months ago

Selected Answer: D

AnswerD

Taking into consideration that we are talking about social media company probably storing a lot of quite small files i would say it cannot be Option A or B.
For example
Amazon S3 File Gateway pricing
Storage Pricing
Storage Pricing Stored and billed as Amazon S3 objects.
Request Pricing
Data written to AWS storage by your gateway $0.01 per GB†
File storage in S3 Billed as Amazon S3 requests.
It can be quite expensive, especially when we will be working on small files.
I would go with EFS with OneZone-IA (option D) which should be less expensive taking into consideration that we are paying only for Storage and Data Transfer (per GB).
But to be honest we need more information to device which solution will be better.
upvoted 1 times

EdricHoang 11 months, 2 weeks ago

It say migrating data into cloud, EFS does not satisfy that. You're right, It needs more information to make your choice to be a better answer.
upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: B

Selected Answer: B

yeah B

upvoted 2 times

seetpt 1 year, 3 months ago

Selected Answer: B

I think B too

upvoted 3 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: B

B looks correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #807 *Topic 1*

A company uses high concurrency AWS Lambda functions to process a constantly increasing number of messages in a message queue during marketing events. The Lambda functions use CPU intensive code to process the messages. The company wants to reduce the compute costs and to maintain service latency for its customers.

Which solution will meet these requirements?

A. Configure reserved concurrency for the Lambda functions. Decrease the memory allocated to the Lambda functions.

B. Configure reserved concurrency for the Lambda functions. Increase the memory according to AWS Compute Optimizer recommendations.

C. Configure provisioned concurrency for the Lambda functions. Decrease the memory allocated to the Lambda functions.

D. Configure provisioned concurrency for the Lambda functions. Increase the memory according to AWS Compute Optimizer recommendations. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (54%) B (41%) A (5%)

Comments

JonJon03 Highly Voted 11 months, 3 weeks ago

Selected Answer: B

eserved concurrency — It guarantees the maximum number of concurrent instances for the function which can be invoked. When a function has being with a reserved concurrency configuration then no other lambda function within the same AWS account and region can use that concurrency. There is no charge for configuring reserved concurrency for a function. Provisioned concurrency — This concurrency initializes a requested number of execution environments so that they are prepared to respond immediately to your function's invocations. Note that configuring provisioned concurrency incurs charges to your AWS account.

upvoted 9 times

tch Most Recent 3 months ago

Selected Answer: B

• Reserved concurrency – This represents the maximum number of concurrent instances allocated to your function. When a function has reserved concurrency, no other function can use that concurrency. Reserved concurrency is useful for ensuring that your most critical functions always have enough concurrency to handle incoming requests. Configuring reserved concurrency for

a function incurs no additional charges.
upvoted 1 times

Gorkha 4 months, 1 week ago

Selected Answer: D

for reserved we don't know the incoming traffic as traffic is increasing so we choose provisioned and as it realted to cpu so compute optimizer.
upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

Selected Answer: D

I vote D instead of B because of constantly increasing number of messages and to maintain service latency for its customer although provisioned concurrency incurs the charges to your AWS account. If the reduce of comput costs is more important than service latency , B is correct.
upvoted 1 times

Salilgen 5 months ago

Selected Answer: D

IMO answer is D.
The number of messages in message queue is constantly increasing: then, to maintain service latency you need to configure provisioned concurrency.
For cpu-intensive job you can optimize the duration increasing memory
upvoted 2 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: B

Let's think - Should we increase or decrease memory? Know that CPU resources scale proportionally with memory allocation in AWS Lambda, so increasing memory will help reduce the execution time for CPU-intensive tasks. Shorter execution times lower the overall compute cost. Then, A & C are out.

About option D, it ensures performance and prevents cold starts, yes, but it also incurs additional costs to keep Lambda execution environments pre-warmed, even when not in use. Reserved concurrency is more cost-effective for sustained workloads.
upvoted 3 times

dragossky 5 months, 3 weeks ago

Selected Answer: D

Reserved concurrency – This represents the maximum number of concurrent instances allocated to your function. When a function has reserved concurrency, no other function can use that concurrency. Reserved concurrency is useful for ensuring that your most critical functions always have enough concurrency to handle incoming requests. Configuring reserved concurrency for a function incurs no additional charges.

Provisioned concurrency – This is the number of pre-initialized execution environments allocated to your function. These execution environments are ready to respond immediately to incoming function requests. Provisioned concurrency is useful for reducing cold start latencies for functions. Configuring provisioned concurrency incurs additional charges to your AWS account.
upvoted 2 times

tch 2 months, 4 weeks ago

then should use reserved concurrency... as no additional charges....
upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: D

Provisioned concurrency:
For scenarios with predictable high traffic like marketing events, using provisioned concurrency ensures Lambda functions are always ready to process requests, maintaining low latency.

Increase memory based on Compute Optimizer recommendations:
Since the Lambda functions are CPU intensive, increasing allocated memory based on AWS Compute Optimizer suggestions can often improve performance and reduce costs by allowing the functions to process more data without hitting CPU throttling.
upvoted 1 times

AMEJack 6 months, 2 weeks ago

Selected Answer: B

Provioned concurrency will increase cost and we need to maintain latency not to decrease latency. Also, provisioned concurrency cost is very related to memory, so if you use provisioned and increase memory this will affect the cost dramatically.

upvoted 2 times

bujuman 1 year ago

Selected Answer: D

Also Lambda provisioned concurrency incure additionnal Account charges (https://docs.aws.amazon.com/lambda/latest/dg/provisioned-concurrency.html), it's the best option because it is stated:
- The company wants to reduce the compute costs and to maintain service latency for its customers.
So maintaining service latency while reducing compute cost is requested.
That being said, Lambda optimization is not a trivial task, that's why one should rely on AWS Compute Optimizer recommendations to analyze usage and find the best fit.
Please read following for more insights:
https://aws.amazon.com/blogs/compute/optimizing-aws-lambda-cost-and-performance-using-aws-compute-optimizer/

upvoted 3 times

JackyCCK 1 year, 2 months ago

Increase the memory according to AWS Compute Optimizer recommendations --> so we can lower the duration of lambda function to reduce the cost.
The ans must be between B & D

upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: D

Provisioned Concurrency keeps the Lambda functions initialized and ready to process incoming events, reducing the cold start latency associated with spinning up new execution environments.

upvoted 4 times

asdfcdsxdfc 1 year, 2 months ago

Selected Answer: D

D is correct

upvoted 3 times

osmk 1 year, 3 months ago

Selected Answer: A

When a large number of messages are in the SQS queue, Lambda scales out, adding additional functions to process the messages. The scale out can consume the concurrency quota in the account. To prevent this from happening, you can set reserved concurrency for individual Lambda functions. This ensures that the specified Lambda function can always scale to that much concurrency, but it also cannot exceed this number.
https://docs.aws.amazon.com/lambda/latest/operatorguide/computing-power.html

upvoted 2 times

osmk 1 year, 3 months ago

When a large number of messages are in the SQS queue, Lambda scales out, adding additional functions to process the messages. The scale out can consume the concurrency quota in the account. To prevent this from happening, you can set reserved concurrency for individual Lambda functions. This ensures that the specified Lambda function can always scale to that much concurrency, but it also cannot exceed this number.
https://docs.aws.amazon.com/lambda/latest/operatorguide/computing-power.html

upvoted 1 times

Sivaeas 1 year, 3 months ago

Selected Answer: D

To reduce compute costs and maintain service latency for customers while using AWS Lambda functions for processing CPU-intensive tasks, you can consider the following strategies:

Optimize Lambda Function Configuration:
Adjust the memory allocation for Lambda functions to better match the CPU requirements of your workload. Higher memory configurations provide more CPU power.
Tune the timeout settings to match the expected processing time of your workload. This prevents unnecessary over-

provisioning and reduces costs.
Fine-tune the concurrency settings to control the number of concurrent executions based on your workload's characteristics.
Use Provisioned Concurrency:
AWS Lambda's provisioned concurrency feature allows you to preallocate a number of execution environments to handle incoming requests instantly. This can help reduce cold starts and maintain consistent performance, especially during peak events.

upvoted 3 times

1dd 1 year, 3 months ago

Reserved concurrency its no charges reduce the computation cost, "latency for its customer" then I'll go for A

upvoted 1 times

lenotc 1 year, 3 months ago

Reserved concurrency guarantees a minimum number of concurrent executions but doesn't inherently improve cold start times like provisioned concurrency.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #808 *Topic 1*

A company runs its workloads on Amazon Elastic Container Service (Amazon ECS). The container images that the ECS task definition uses need to be scanned for Common Vulnerabilities and Exposures (CVEs). New container images that are created also need to be scanned.

Which solution will meet these requirements with the FEWEST changes to the workloads?

A. Use Amazon Elastic Container Registry (Amazon ECR) as a private image repository to store the container images. Specify scan on push filters for the ECR basic scan. **Most Voted**

B. Store the container images in an Amazon S3 bucket. Use Amazon Macie to scan the images. Use an S3 Event Notification to initiate a Macie scan for every event with an s3:ObjectCreated:Put event type.

C. Deploy the workloads to Amazon Elastic Kubernetes Service (Amazon EKS). Use Amazon Elastic Container Registry (Amazon ECR) as a private image repository. Specify scan on push filters for the ECR enhanced scan.

D. Store the container images in an Amazon S3 bucket that has versioning enabled. Configure an S3 Event Notification for s3:ObjectCreated:* events to invoke an AWS Lambda function. Configure the Lambda function to initiate an Amazon Inspector scan.

Correct Answer: *A*

Community vote distribution

A (95%) D (5%)

Comments

1dd Highly Voted 9 months ago

Selected Answer: A

need less workload changes and CVEs
https://docs.aws.amazon.com/AmazonECR/latest/userguide/image-scanning.html

upvoted 9 times

xBUGx Highly Voted 9 months ago

Selected Answer: A

FEWEST changes to the workloads and scan CVE is enough. A looks OK.

upvoted 6 times

zdi561 Most Recent 3 months, 3 weeks ago

Selected Answer: D

A is not right because ECR basic scan only scan OS. in B enhanced scan does scan app packages but it is required to run EKS. D is simpler.

upvoted 1 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: A

Amazon ECR: Designed for storing container images, integrates seamlessly with Amazon ECS.
Scan on Push: Automatically scans images for vulnerabilities (CVEs) when they are uploaded. This is built-in functionality and requires minimal changes to existing workflows.
Fewest Changes: Existing ECS workloads can directly pull images from ECR without modification to the workloads themselves.

B - Amazon Macie focuses on data discovery and classification (such as identifying sensitive data in S3 buckets), not container image vulnerabilities.
C - Result-wise OK, but switching from ECS to EKS is unnecessary and involves a steep learning curve and operational overhead.
D - Result-wise OK. Amazon Inspector can scan for vulnerabilities, but this is too complex and requires a custom setup, including using S3 and Lambda to trigger scans.

upvoted 1 times

sandordini 7 months, 2 weeks ago

Selected Answer: A

Basic scan looks for Common Vulnerabilities and Exposures (CVEs)

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #809 *Topic 1*

A company uses an AWS Batch job to run its end-of-day sales process. The company needs a serverless solution that will invoke a third-party reporting application when the AWS Batch job is successful. The reporting application has an HTTP API interface that uses username and password authentication.

Which solution will meet these requirements?

A. Configure an Amazon EventBridge rule to match incoming AWS Batch job SUCCEEDED events. Configure the third-party API as an EventBridge API destination with a username and password. Set the API destination as the EventBridge rule target. **Most Voted**

B. Configure Amazon EventBridge Scheduler to match incoming AWS Batch job SUCCEEDED events. Configure an AWS Lambda function to invoke the third-party API by using a username and password. Set the Lambda function as the EventBridge rule target.

C. Configure an AWS Batch job to publish job SUCCEEDED events to an Amazon API Gateway REST API. Configure an HTTP proxy integration on the API Gateway REST API to invoke the third-party API by using a username and password.

D. Configure an AWS Batch job to publish job SUCCEEDED events to an Amazon API Gateway REST API. Configure a proxy integration on the API Gateway REST API to an AWS Lambda function. Configure the Lambda function to invoke the third-party API by using a username and password.

Correct Answer: *A*

Community vote distribution

A (65%) B (26%) D (9%)

Comments

venutadi Highly Voted 1 year, 1 month ago

Selected Answer: A

https://aws.amazon.com/blogs/compute/using-api-destinations-with-amazon-eventbridge/
Amazon EventBridge enables developers to route events between AWS services, integrated software as a service (SaaS) applications, and your own applications. It can help decouple applications and produce more extensible, maintainable architectures. With the new API destinations feature, EventBridge can now integrate with services outside of AWS using REST API calls.

upvoted 13 times

shintaro0914 1 year, 1 month ago

I agree.
upvoted 1 times

sandordini Highly Voted 1 year, 1 month ago

I'm confused. Both A and B seem to be viable. There is no requirement of cost, complexity, or overhead. :S
upvoted 5 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: A

A - EventBridge API Destination is specifically designed to send events to HTTP APIs. It supports basic authentication (username and password), which fits the requirement for 3rd-party API authentication.
B - OK, but EventBridge Scheduler is unnecessary because EventBridge rules can already match Batch job events. Introducing Lambda functions adds complexity and costs.
C - An HTTP proxy integration on API Gateway cannot handle authentication directly. We still need to manage username and password authentication separately.
D - OK, but even more complex than option B.
upvoted 3 times

Scheldon 12 months ago

Selected Answer: A

AnswerA

EventBridge Schedule will not work as it will allow us to "do something" per schedule.
EventBridge rule will allow us to "do something" when event will occur.
I think there is no possibility to publish/send job "SUCCEEDED" to AMAZON API Gateway REST API or that we can do anykind of integration with AMAZON API Gateway, hence I would choose A
upvoted 1 times

bujuman 1 year ago

Selected Answer: A

Even though option A and B could do the trick and also no statement related to least effort is rquested, EventBridge is dedicated for similar use case.
https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-api-destinations.html
Plus, it can also handle basic authentication
https://aws.amazon.com/blogs/compute/using-api-destinations-with-amazon-eventbridge/
upvoted 1 times

Sergiuss95 1 year, 1 month ago

Selected Answer: B

I think is better to programming a lambda and obtain user and password from Secret Manager... So I think the better solution is B
upvoted 3 times

EdricHoang 11 months, 1 week ago

EventBridge also store credentials in Secret manager: https://aws.amazon.com/blogs/compute/using-api-destinations-with-amazon-eventbridge/
upvoted 1 times

Oluwatosin09 1 year, 1 month ago

Selected Answer: B

Answer should be B.
upvoted 1 times

boluwatito 1 year, 1 month ago

Selected Answer: B

Configure Amazon EventBridge Scheduler to match incoming AWS Batch job SUCCEEDED events.
Configure an AWS Lambda function to invoke the third-party API using a username and password.
Set the Lambda function as the EventBridge rule target.

upvoted 2 times

AlvinC2024 1 year, 2 months ago

Selected Answer: A

A. Configure an Amazon EventBridge rule to match incoming AWS Batch job SUCCEEDED events. Configure the third-party API as an EventBridge API destination with a username and password. Set the API destination as the EventBridge rule target.

This option is the most direct and serverless approach to meeting the requirements. Amazon EventBridge can detect the successful completion of the AWS Batch job and trigger actions based on this event. By configuring the third-party API as an API destination with authentication credentials, EventBridge can directly invoke the third-party reporting application without the need for additional services. This approach minimizes complexity and operational overhead.

upvoted 4 times

alawada 1 year, 2 months ago

Selected Answer: D

Create an AWS Lambda function responsible for invoking the third-party reporting application's HTTP API endpoint. The Lambda function will be triggered by the successful completion of the AWS Batch job.

upvoted 3 times

k_k_kkk 1 year, 2 months ago

Selected Answer: B

AWS Batch sends job status change to EventBridge.

https://docs.aws.amazon.com/batch/latest/userguide/batch_cwe_events.html

upvoted 3 times

osmk 1 year, 3 months ago

look like B

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #810 *Topic 1*

A company collects and processes data from a vendor. The vendor stores its data in an Amazon RDS for MySQL database in the vendor's own AWS account. The company's VPC does not have an internet gateway, an AWS Direct Connect connection, or an AWS Site-to-Site VPN connection. The company needs to access the data that is in the vendor database.

Which solution will meet this requirement?

A. Instruct the vendor to sign up for the AWS Hosted Connection Direct Connect Program. Use VPC peering to connect the company's VPC and the vendor's VPC.

B. Configure a client VPN connection between the company's VPC and the vendor's VPC. Use VPC peering to connect the company's VPC and the vendor's VPC.

C. Instruct the vendor to create a Network Load Balancer (NLB). Place the NLB in front of the Amazon RDS for MySQL database. Use AWS PrivateLink to integrate the company's VPC and the vendor's VPC. **Most Voted**

D. Use AWS Transit Gateway to integrate the company's VPC and the vendor's VPC. Use VPC peering to connect the company's VPC and the vendor's VPC.

Correct Answer: *C*

Community vote distribution

C (95%) A (5%)

Comments

Ucy Highly Voted 10 months, 2 weeks ago

Pour yourself a cold beer, when you get to this question, its been a very long run

upvoted 16 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: C

PrivateLink: AWS PrivateLink enables secure and private access to resources (like RDS) across accounts and VPCs using interface endpoints. It avoids exposing the database to the internet.
NLB: A Network Load Balancer acts as the integration point for PrivateLink. The vendor's NLB will route traffic to the RDS database.
VPC Without Internet Gateway: PrivateLink works even if the company's VPC does not have an internet gateway because communication occurs over private AWS infrastructure.

Option A, B, and D all mention VPC Peering, but it doesn't enable direct access to RDS without additional routing or security configurations.

upvoted 4 times

Scheldon 12 months ago

Selected Answer: C

AnswerC

AWS PrivateLink enables you to connect to some AWS services, services hosted by other AWS accounts (referred to as endpoint services), and supported AWS Marketplace partner services, via private IP addresses in your VPC. The interface endpoints are created directly inside of your VPC, using elastic network interfaces and IP addresses in your VPC’s subnets. That means that VPC Security Groups can be used to manage access to the endpoints.
https://docs.aws.amazon.com/whitepapers/latest/aws-vpc-connectivity-options/aws-privatelink.html

upvoted 2 times

Nm55569 1 year ago

Selected Answer: C

https://aws.amazon.com/blogs/database/access-amazon-rds-across-vpcs-using-aws-privatelink-and-network-load-balancer/

upvoted 2 times

TwinSpark 1 year ago

I think i go for C, because if you exclude Dirct connect, VPN and GW so only C is available. but create an NLB zo do not want provision a transit GW sounds weird to me

upvoted 1 times

sandordini 1 year, 1 month ago

Selected Answer: C

Private link:
Does not require VPC linking: NO Internet Gateway, NO NAT GAteway, No Route table
Needs NLB on Service VPC, and ENI on the Customer VPC

upvoted 4 times

rondelldell 1 year, 2 months ago

D
You can peer both intra-Region and inter-Region transit gateways, and route traffic between them, which includes IPv4 and IPv6 traffic. To do this, create a peering attachment on your transit gateway, and specify a transit gateway. The peer transit gateway can be in your account or a different AWS account.

After you create a peering attachment request, the owner of the peer transit gateway (also referred to as the accepter transit gateway) must accept the request. To route traffic between the transit gateways, add a static route to the transit gateway route table that points to the transit gateway peering attachment.
https://docs.aws.amazon.com/vpc/latest/tgw/tgw-peering.html

upvoted 2 times

xBUGx 1 year, 2 months ago

D does not involve internet. But TGW is unnecessary.
A is more simple and clear.

upvoted 1 times

Sivaeas 1 year, 3 months ago

Selected Answer: C

AWS PrivateLink:
AWS PrivateLink enables you to privately access services hosted on AWS in a highly available and scalable manner. With PrivateLink, you can access the vendor's RDS for MySQL instance securely without exposing it to the public internet.
The vendor can create a VPC endpoint for RDS within their own VPC, which acts as an entry point for accessing the RDS instance. This endpoint can then be shared with the company.
The company can create a VPC endpoint service in their VPC and accept the endpoint connection request from the vendor. This allows the company's resources to communicate with the RDS instance securely through PrivateLink.

upvoted 3 times

lenotc 1 year, 3 months ago

Selected Answer: C

C is correct:
https://aws.amazon.com/blogs/networking-and-content-delivery/how-to-securely-publish-internet-applications-at-scale-using-application-load-balancer-and-aws-privatelink/
upvoted 1 times

1dd 1 year, 3 months ago

Selected Answer: C

Plz commit the previous comment,
A involve- Direct connect
B involve - peering required same region
D involve - uses internet gateway
upvoted 3 times

1dd 1 year, 3 months ago

Selected Answer: A

No internet gateway XD
No Direct connect XC
No Peering XB
upvoted 1 times

asdfcdsxdfc 1 year, 3 months ago

Shouldn't it be D?
upvoted 3 times

Sergantus 6 months, 4 weeks ago

VPC peering merges two VPCs and exposes all the services across both VPCs, which is more than less desirable
upvoted 1 times

rondelldell 1 year, 2 months ago

YES D
transit gateway is like router - u can connect VPCs AND OnPrem. VPCs can be in anuther account or region or org
upvoted 1 times

1dd 1 year, 3 months ago

I think it required use of internet gateway .
upvoted 1 times

Jacky_S 11 months, 2 weeks ago

No, it did not
https://docs.aws.amazon.com/vpc/latest/tgw/tgw-transit-gateways.html
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #811 *Topic 1*

A company wants to set up Amazon Managed Grafana as its visualization tool. The company wants to visualize data from its Amazon RDS database as one data source. The company needs a secure solution that will not expose the data over the internet.

Which solution will meet these requirements?

A. Create an Amazon Managed Grafana workspace without a VPC. Create a public endpoint for the RDS database. Configure the public endpoint as a data source in Amazon Managed Grafana.

B. Create an Amazon Managed Grafana workspace in a VPC. Create a private endpoint for the RDS database. Configure the private endpoint as a data source in Amazon Managed Grafana. Most Voted

C. Create an Amazon Managed Grafana workspace without a VPCCreate an AWS PrivateLink endpoint to establish a connection between Amazon Managed Grafana and Amazon RDS. Set up Amazon RDS as a data source in Amazon Managed Grafana.

D. Create an Amazon Managed Grafana workspace in a VPC. Create a public endpoint for the RDS database. Configure the public endpoint as a data source in Amazon Managed Grafana.

Correct Answer: *B*

Community vote distribution

B (66%) C (34%)

Comments

Sergiuss95 Highly Voted 1 year, 1 month ago

Selected Answer: B

I think is b. Private endpoint sounds like private vpc endpoint, that is equals to privatelink

upvoted 6 times

Bazzix Highly Voted 1 year, 2 months ago

Selected Answer: B

B is correct

upvoted 5 times

tch Most Recent 2 months, 4 weeks ago

Selected Answer: C

should be C:
https://docs.aws.amazon.com/grafana/latest/userguide/AMG-configure-nac.html
Amazon Managed Grafana workspace not created in VPC.....
upvoted 1 times

GOTJ 3 months, 1 week ago

Selected Answer: B

With all due respect, once again options are poorly written: you don't create an Amazon Managed Grafana workspace IN or WITHOUT a VPC. Instead, you create an AMG WITH or WITHOUT a VPC CONNECTION:

https://docs.aws.amazon.com/grafana/latest/userguide/AMG-create-workspace.html
(Step 8, Optional): You can choose to connect to an Amazon virtual private cloud (VPC) on this page, or you can connect to a VPC later. To learn more, see Connect to data sources or notification channels in Amazon VPC from Amazon Managed Grafana.

According to the "to learn more" link, if you create the AMG WITHOUT a VPC CONNECTION:

By default, traffic from your Amazon Managed Grafana workspace to data sources or notification channels flows via the public Internet. This limits the connectivity from your Amazon Managed Grafana workspace to services that are publicly accessible.

So option "C" must be discarded as it doesn't comply company requirements of avoid using internet
upvoted 1 times

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: B

A & D are out, because they're using public endpoints, which causes exposure over the Internet.

Why C is wrong: If the Grafana workspace is not in a VPC, it cannot use PrivateLink endpoints. PrivateLink endpoints must reside within a VPC. If Grafana is outside the VPC, PrivateLink cannot establish the required secure, private connection.
upvoted 4 times

ARV14 6 months, 1 week ago

Selected Answer: C

https://docs.aws.amazon.com/grafana/latest/userguide/VPC-endpoints.html , https://aws.amazon.com/about-aws/whats-new/2022/11/amazon-managed-grafana-connection-data-sources-hosted-virtual-private-cloud/
upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

By default, traffic from your Amazon Managed Grafana workspace to data sources or notification channels flows via the public Internet. This limits the connectivity from your Amazon Managed Grafana workspace to services that are publicly accessible.
If you want to connect to private-facing data sources that are within a VPC, or keep traffic local to a VPC, you can connect your Amazon Managed Grafana workspace to the Amazon Virtual Private Cloud (Amazon VPC) hosting these data sources. After you configure the VPC data source connection, all traffic flows via your VPC.
upvoted 1 times

chest_jd 6 months, 4 weeks ago

Selected Answer: B

Choice B or C could be resolved in this way:
B. Create an Amazon Managed Grafana workspace in a VPC
C. Create an Amazon Managed Grafana workspace without a VPC

As far as I know we cannot create workspace without VCP
upvoted 3 times

tonybuivannghia 7 months, 2 weeks ago

Selected Answer: C

After searching effort, I agree C is correct because AMG workspace can't include in VPC. When you have not configured a private VPC, and Amazon Managed Grafana is connecting to publicly accessible data sources, it connects to some AWS services in the same region via AWS PrivateLink. This includes services such as CloudWatch, Amazon Managed Service for Prometheus and AWS X-Ray. Traffic to those services does not flow via the public Internet.
upvoted 1 times

NSA_Poker 10 months, 2 weeks ago

Selected Answer: C

(B or C)?-1 = Do we create AMG workspace in a VPC OR do we create AMG workspace without a VPC? AMG is NOT created within a VPC; AMG connects to a VPC. "Currently, you can connect one Amazon Managed Grafana workspace to one VPC endpoint in the same region and same account. However, you can use Virtual Private Cloud peering or AWS Transit Gateway to connect the cross-region or cross-account VPCs, then connect the select the VPC endpoint that's in the same account and same region as your Amazon Managed Grafana workspace." -FAQs

(C) is correct.

upvoted 3 times

NSA_Poker 10 months, 2 weeks ago

(B or C)?-2 = private endpoint OR AWS PrivateLink? The brand-name is more correct.

(B or C)?-3 = Configure the private endpoint as a data source OR Set up Amazon RDS as a data source? In the AMG console, after clicking on Data sources, you'll see a list of AWS services (Athena, Redshift etc) NOT network endpoints. After selecting RDS, you can further specify the Region & Resource ID.

(B) eliminated.
(C) is correct.

upvoted 1 times

LeonSauveterre 5 months, 2 weeks ago

You're missing out PrivateLink in option C.

upvoted 1 times

EdricHoang 11 months, 3 weeks ago

Selected Answer: B

Its B.
C is also a valid choice
"Not exposing to the internet" is letting me eliminate C

upvoted 3 times

NSA_Poker 11 months ago

(B) "a private endpoint" & (C) "an AWS PrivateLink endpoint" do NOT expose traffic to the internet.
(A & D) eliminated. "a public endpoint for the RDS database" would "expose the data over the internet"

upvoted 1 times

ike001 11 months, 3 weeks ago

B as you need to create Managed Grafana workspace with a VPC for private access
https://docs.aws.amazon.com/grafana/latest/userguide/AMG-configure-nac.html

upvoted 2 times

NSA_Poker 10 months, 1 week ago

(B) doesn't say 'with a VPC'; it says "..Grafana workspace IN A VPC."
(B) eliminated.
(C) is correct.

upvoted 2 times

Nm55569 1 year ago

Selected Answer: B

https://aws.amazon.com/about-aws/whats-new/2022/11/amazon-managed-grafana-connection-data-sources-hosted-virtual-private-cloud/

upvoted 3 times

sandordini 1 year, 1 month ago

I guess they mean C, But again, it's strange... IMO B would also work... There is no requirement for the least effort... Pls, correct me if I'm wrong...

upvoted 2 times

venutadi 1 year, 1 month ago

Selected Answer: C

Once you configure direct connectivity between a Grafana workspace and a VPC, Amazon Managed Grafana creates and manages an elastic network interface (ENI) per subnet to connect to the VPC. This enables the Grafana workspace to connect to data sources within the VPC, such as OpenSearch domains or RDS databases. Additionally, all traffic is now routed through the configured VPC, including alert destination and data source connectivity.

upvoted 4 times

VortexMD 1 year, 2 months ago

AWS PrivateLink provides private connectivity between virtual private clouds (VPCs), supported AWS services, and your on-premises networks without exposing your traffic to the public internet. Interface VPC endpoints, powered by PrivateLink, connect you to services hosted by AWS Partners and supported solutions available in AWS Marketplace.

upvoted 1 times

VortexMD 1 year, 2 months ago

https://aws.amazon.com/blogs/mt/announcing-private-vpc-data-source-support-for-amazon-managed-grafana/

upvoted 1 times

osmk 1 year, 2 months ago

Selected Answer: C

cccc ccc

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #812 *Topic 1*

A company hosts a data lake on Amazon S3. The data lake ingests data in Apache Parquet format from various data sources. The company uses multiple transformation steps to prepare the ingested data. The steps include filtering of anomalies, normalizing of data to standard date and time values, and generation of aggregates for analyses.

The company must store the transformed data in S3 buckets that data analysts access. The company needs a prebuilt solution for data transformation that does not require code. The solution must provide data lineage and data profiling. The company needs to share the data transformation steps with employees throughout the company.

Which solution will meet these requirements?

A. Configure an AWS Glue Studio visual canvas to transform the data. Share the transformation steps with employees by using AWS Glue jobs.

B. Configure Amazon EMR Serverless to transform the data. Share the transformation steps with employees by using EMR Serverless jobs.

C. Configure AWS Glue DataBrew to transform the data. Share the transformation steps with employees by using DataBrew recipes. **Most Voted**

D. Create Amazon Athena tables for the data. Write Athena SQL queries to transform the data. Share the Athena SQL queries with employees.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: C

A - AWS Glue doesn't focus on data profiling or provide an easy way for non-technical employees to collaborate on transformation steps.
B - You need to code.
C - AWS Glue DataBrew is a no-code tool specifically designed for data preparation and transformation.
D - Let's say writing SQLs are not coding, still, Athena lacks data lineage, data profiling, and the ability to reuse transformations as recipes.

upvoted 2 times

Lingo43 11 months, 1 week ago

Selected Answer: C

AWS Glue DataBrew: This is a visual data preparation tool that allows you to clean and normalize data without writing code. It has built-in transformations for common tasks like filtering anomalies, normalizing dates, and generating aggregates. It also provides data lineage and profiling capabilities, which are required by the company.
DataBrew Recipes: These are reusable workflows that define the data transformation steps. They can be easily shared with other employees, making it simple to collaborate on data preparation tasks.

upvoted 2 times

Scheldon 12 months ago

Selected Answer: C

AnswerC

AWS Glue DataBrew is a visual data preparation tool that enables users to clean and normalize data without writing any code. Using DataBrew helps reduce the time it takes to prepare data for analytics and machine learning (ML) by up to 80 percent, compared to custom developed data preparation. You can choose from over 250 ready-made transformations to automate data preparation tasks, such as filtering anomalies, converting data to standard formats, and correcting invalid values.
https://docs.aws.amazon.com/databrew/latest/dg/what-is.html

upvoted 4 times

Linuslin 1 year ago

Selected Answer: C

C is correct.
https://docs.aws.amazon.com/databrew/latest/dg/recipes.html

upvoted 1 times

seetpt 1 year, 3 months ago

Selected Answer: C

Agree with C

upvoted 2 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: C

Should be C

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #813 *Topic 1*

A solutions architect runs a web application on multiple Amazon EC2 instances that are in individual target groups behind an Application Load Balancer (ALB). Users can reach the application through a public website.

The solutions architect wants to allow engineers to use a development version of the website to access one specific development EC2 instance to test new features for the application. The solutions architect wants to use an Amazon Route 53 hosted zone to give the engineers access to the development instance. The solution must automatically route to the development instance even if the development instance is replaced.

Which solution will meet these requirements?

A. Create an A Record for the development website that has the value set to the ALB. Create a listener rule on the ALB that forwards requests for the development website to the target group that contains the development instance. Most Voted

B. Recreate the development instance with a public IP address. Create an A Record for the development website that has the value set to the public IP address of the development instance.

C. Create an A Record for the development website that has the value set to the ALB. Create a listener rule on the ALB to redirect requests for the development website to the public IP address of the development instance.

D. Place all the instances in the same target group. Create an A Record for the development website. Set the value to the ALB. Create a listener rule on the ALB that forwards requests for the development website to the target group.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

Mikado211 Highly Voted 1 year, 2 months ago

Both A and C look correct but with the C you pass through the ALB to be redirected to a public IP (so go outside) to come back again through this public IP which is not ideal.

The answer A is much cleaner and simplier with a dedicated target group and a listener rule pointing it.

upvoted 7 times

MatAlves 8 months, 3 weeks ago

B: Directly using a public IP address ties the A Record to a specific instance, which is not ideal since replacing the instance would require manually updating the Route 53 record.

C: Redirecting to a public IP address of the development instance is similar to option B, and would also require manual updates if the instance changes.
upvoted 2 times

elmyth 9 months ago

Public IP is not permanent, so definitely not B and C
upvoted 3 times

gdf54634 Highly Voted 1 year, 2 months ago

Selected Answer: A

Should be A as it points to the target group for easy replacement etc
upvoted 7 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: A

B - If the development instance is replaced, the public IP address changes, requiring manual DNS updates in Route 53.
C - The ALB listener rule cannot directly redirect traffic based on the instance's public IP address without using a target group.
D - Without traffic isolation, you could easily route requests for the development website to production instances. Dangerous and just plain wrong.
upvoted 1 times

LeonSauveterre 5 months, 2 weeks ago

FYI, I copied this after googling.
What is a DNS "A record"?
The "A" stands for "address" and this is the most fundamental type of DNS record: it indicates the IP address of a given domain. For example, if you pull the DNS records of cloudflare.com, the A record currently returns an IP address of: 104.17.210.9.

A records only hold IPv4 addresses. If a website has an IPv6 address, it will instead use an "AAAA" record.
upvoted 1 times

Cpso 6 months, 1 week ago

Selected Answer: A

A.But not totally correct. ALB IP can be change. should use CNAME record.
upvoted 2 times

GOTJ 2 months, 2 weeks ago

That's right! In fact, you shouldn't associate an "A" record to an ALB. You must use a CNAME record instead
upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: A

By setting the A Record to point to the ALB, you ensure that even if the development instance is replaced, the DNS record will still direct traffic to the correct target group which contains the new development instance, maintaining accessibility for engineers.
upvoted 1 times

JA2018 6 months, 1 week ago

Why we have to rule out the other options:
B: Using the public IP address of the development instance directly would break access if the instance is replaced, as the IP address would change.
C: Redirecting requests from the ALB to the development instance's public IP is not a good practice as it bypasses the load balancing functionality of the ALB.
D: Placing all instances in the same target group would not allow for separate access to the development instance as the load balancer would distribute traffic across all instances, including production ones.
upvoted 1 times

JA2018 6 months, 1 week ago

To conclude, option A provides the necessary flexibility and reliability to access the development instance while utilizing the ALB for load balancing and automatic failover capabilities.

upvoted 1 times

Scheldon 12 months ago

Selected Answer: A

AnswerA,

With ALB which will point to appopriate dev group we will be able easy to create HA for dev servers.

upvoted 3 times

asdfcdsxdfc 1 year, 2 months ago

Selected Answer: A

I think its A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #814 | *Topic 1*

A company runs a container application on a Kubernetes cluster in the company's data center. The application uses Advanced Message Queuing Protocol (AMQP) to communicate with a message queue. The data center cannot scale fast enough to meet the company's expanding business needs. The company wants to migrate the workloads to AWS.

Which solution will meet these requirements with the LEAST operational overhead?

A. Migrate the container application to Amazon Elastic Container Service (Amazon ECS). Use Amazon Simple Queue Service (Amazon SQS) to retrieve the messages.

B. Migrate the container application to Amazon Elastic Kubernetes Service (Amazon EKS). Use Amazon MQ to retrieve the messages. **Most Voted**

C. Use highly available Amazon EC2 instances to run the application. Use Amazon MQ to retrieve the messages.

D. Use AWS Lambda functions to run the application. Use Amazon Simple Queue Service (Amazon SQS) to retrieve the messages.

Correct Answer: *B*

Community vote distribution

B (95%) A (5%)

Comments

Mikado211 Highly Voted 1 year, 2 months ago

Selected Answer: B

This question is a trap because A is definitely the answer for a Least overhead (ECS + SQS) and in a real life scenario could be good in 99% of cases.

However SQS do not implement AMQP (SQS is only a simple queueing system very basic) so we have to use Amazon MQ.

In terms of containers EKS will always be a better solution than a manual setup of Docker.

Good solution would have been ECS+AmazonMQ not given here

Lambda can work with containers, but since there are limitations like 15 minutes limit we can't really consider it as a good solution.

So B is the least bad solution.

upvoted 15 times

Dantecito Most Recent 3 months, 2 weeks ago

Selected Answer: B

We were asked for the LEAST operational overhead, so the company already has a managed message broker and thats why we use Amazon MQ and they are already using Kubernetes.

upvoted 1 times

Scheldon 12 months ago

Selected Answer: B

AnswerB

For me B is a correct solution. In question AMQP is mentioned and Amazon on his doc page about MQ is providing such information:
"Amazon MQ is a managed message broker service that provides compatibility with many popular message brokers. We recommend Amazon MQ for migrating applications from existing message brokers that rely on compatibility with APIs such as JMS or protocols such as AMQP 0-9-1, AMQP 1.0, MQTT, OpenWire, and STOMP."
I cannot be coincidence that documentation is mentioning about AMQP.

https://docs.aws.amazon.com/amazon-mq/latest/developer-guide/welcome.html

upvoted 1 times

sandordini 1 year, 1 month ago

Selected Answer: A

I'd go for A.
Although the ideal solution and least modification would require B, with heavy rework the application can be (most likely) adopted to ECS+SQS. As it is an AWS exam, not a vendor-agnostic SA exam, A will be the correct answer.

upvoted 1 times

seetpt 1 year, 3 months ago

Selected Answer: B

B because only solution with Kubernetes

upvoted 1 times

asdfcdsxdfc 1 year, 3 months ago

Selected Answer: B

Should be B

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #815 *Topic 1*

An online gaming company hosts its platform on Amazon EC2 instances behind Network Load Balancers (NLBs) across multiple AWS Regions. The NLBs can route requests to targets over the internet. The company wants to improve the customer playing experience by reducing end-to-end load time for its global customer base.

Which solution will meet these requirements?

A. Create Application Load Balancers (ALBs) in each Region to replace the existing NLBs. Register the existing EC2 instances as targets for the ALBs in each Region.

B. Configure Amazon Route 53 to route equally weighted traffic to the NLBs in each Region.

C. Create additional NLBs and EC2 instances in other Regions where the company has large customer bases.

D. Create a standard accelerator in AWS Global Accelerator. Configure the existing NLBs as target endpoints. Most Voted

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

Mikado211 Highly Voted 1 year, 2 months ago

Selected Answer: D

In such situation if you had an ALB you would use Cloudfront
Since you have a NLB you use AWS Global Accelerator
So D.

upvoted 8 times

asdfcdsxdfc Highly Voted 1 year, 3 months ago

Selected Answer: D

Should be D

upvoted 5 times

Scheldon Most Recent 12 months ago

Selected Answer: D

AnswerD

Usage of Global Accelerator should help here.
"
Acceleration for latency-sensitive applications
Many applications, especially in areas such as gaming, media, mobile apps, ad-tech, and financials, require very low latency for a great user experience. To improve the user experience, Global Accelerator directs user traffic to the application endpoint that is nearest to the client, which reduces internet latency and jitter. Global Accelerator routes traffic to the closest edge location by using Anycast, and then routes it to the closest regional endpoint over the AWS global network. Global Accelerator quickly reacts to changes in network performance to improve your users' application performance.
"
https://docs.aws.amazon.com/global-accelerator/latest/dg/introduction-benefits-of-migrating.html

upvoted 3 times

seetpt 1 year, 3 months ago

Selected Answer: D

Agree with D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #816 *Topic 1*

A company has an on-premises application that uses SFTP to collect financial data from multiple vendors. The company is migrating to the AWS Cloud. The company has created an application that uses Amazon S3 APIs to upload files from vendors.

Some vendors run their systems on legacy applications that do not support S3 APIs. The vendors want to continue to use SFTP-based applications to upload data. The company wants to use managed services for the needs of the vendors that use legacy applications.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an AWS Database Migration Service (AWS DMS) instance to replicate data from the storage of the vendors that use legacy applications to Amazon S3. Provide the vendors with the credentials to access the AWS DMS instance.

B. Create an AWS Transfer Family endpoint for vendors that use legacy applications. **Most Voted**

C. Configure an Amazon EC2 instance to run an SFTP server. Instruct the vendors that use legacy applications to use the SFTP server to upload data.

D. Configure an Amazon S3 File Gateway for vendors that use legacy applications to upload files to an SMB file share.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

asdfcdsxdfc Highly Voted 1 year, 3 months ago

Selected Answer: B

B is correct

upvoted 8 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: B

A - DMS is designed for database migrations, not for file-based SFTP uploads.
B - AWS Transfer Family is a fully managed service that supports SFTP, FTPS, and FTP protocols.
C - OK but significant operational overhead.
D - S3 File Gateway supports SMB and NFS, not SFTP.

D - S3 File Gateway supports SMB and NFS, not SFTP.
upvoted 1 times

Dantecito 3 months, 2 weeks ago

also C - Not a managed services.
upvoted 1 times

Johnoppong101 9 months, 3 weeks ago

Selected Answer: B

Answer is B
upvoted 1 times

Sergiuss95 1 year, 1 month ago

Selected Answer: B

Explanation:

AWS Transfer Family is a fully managed service that allows you to set up SFTP, FTPS, and FTP endpoints for accessing Amazon S3 and Amazon EFS storage.
By creating an AWS Transfer Family endpoint, the company can provide vendors with the familiar SFTP interface to upload data directly to Amazon S3 without requiring them to make any changes to their legacy applications.
This solution eliminates the need for the company to manage and maintain additional infrastructure such as EC2 instances or file gateways.
AWS Transfer Family handles scalability, availability, and security, reducing operational overhead for the company.
upvoted 4 times

seetpt 1 year, 3 months ago

Selected Answer: B

B is correct
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #817 *Topic 1*

A marketing team wants to build a campaign for an upcoming multi-sport event. The team has news reports from the past five years in PDF format. The team needs a solution to extract insights about the content and the sentiment of the news reports. The solution must use Amazon Textract to process the news reports.

Which solution will meet these requirements with the LEAST operational overhead?

A. Provide the extracted insights to Amazon Athena for analysis. Store the extracted insights and analysis in an Amazon S3 bucket.

B. Store the extracted insights in an Amazon DynamoDB table. Use Amazon SageMaker to build a sentiment model.

C. Provide the extracted insights to Amazon Comprehend for analysis. Save the analysis to an Amazon S3 bucket.
Most Voted

D. Store the extracted insights in an Amazon S3 bucket. Use Amazon QuickSight to visualize and analyze the data.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

Scheldon 12 months ago

Selected Answer: C

AnswerC
"
Amazon Comprehend uses natural language processing (NLP) to extract insights about the content of documents. It develops insights by recognizing the entities, key phrases, language, sentiments, and other common elements in a document. Use Amazon Comprehend to create new products based on understanding the structure of documents. For example, using Amazon Comprehend you can search social networking feeds for mentions of products or scan an entire document repository for key phrases.
"
https://docs.aws.amazon.com/comprehend/latest/dg/what-is.html

upvoted 3 times

zinabu 1 year, 1 month ago

Selected Answer: C

Selected Answer: C
amazon comprehend= sentiment analysis
upvoted 2 times

zinabu 1 year, 2 months ago

Selected Answer: C
amazon comprehend= sentiment analysis
upvoted 4 times

alawada 1 year, 2 months ago

Selected Answer: C

Whenever new PDF files are uploaded to the designated S3 bucket, the Lambda function will be triggered to extract insights using Textract and Comprehend.
upvoted 3 times

Mikado211 1 year, 2 months ago

Selected Answer: C

When you have words like "sentiment" in a sentence, it's related to Comprehend
So C.
upvoted 2 times

seetpt 1 year, 3 months ago

Selected Answer: C

Maybe C?
upvoted 2 times

asdfcdsxdfc 1 year, 3 months ago

Shouldn't it be C?
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #818 *Topic 1*

A company's application runs on Amazon EC2 instances that are in multiple Availability Zones. The application needs to ingest real-time data from third-party applications.

The company needs a data ingestion solution that places the ingested raw data in an Amazon S3 bucket.

Which solution will meet these requirements?

A. Create Amazon Kinesis data streams for data ingestion. Create Amazon Kinesis Data Firehose delivery streams to consume the Kinesis data streams. Specify the S3 bucket as the destination of the delivery streams. **Most Voted**

B. Create database migration tasks in AWS Database Migration Service (AWS DMS). Specify replication instances of the EC2 instances as the source endpoints. Specify the S3 bucket as the target endpoint. Set the migration type to migrate existing data and replicate ongoing changes.

C. Create and configure AWS DataSync agents on the EC2 instances. Configure DataSync tasks to transfer data from the EC2 instances to the S3 bucket.

D. Create an AWS Direct Connect connection to the application for data ingestion. Create Amazon Kinesis Data Firehose delivery streams to consume direct PUT operations from the application. Specify the S3 bucket as the destination of the delivery streams.

Correct Answer: *A*

Community vote distribution

A (78%) C (22%)

Comments

asdfcdsxdfc Highly Voted 1 year, 3 months ago

Selected Answer: A

A is correct

upvoted 9 times

JA2018 Most Recent 6 months, 1 week ago

Selected Answer: A

Why A?