I would go with Lambda and SQS.
when using SQS we will be sure that all images will be processed and hence to process we need 2 min and 512 MB of memory (Lambda is allowing upto 15 min and upto10K MB) Lambda should be perfect scalable solution which will allow for almost in real time image processing.

upvoted 8 times

Scheldon 1 year ago

and it is cost effective ;)

upvoted 2 times

kevindu Most Recent 9 months, 3 weeks ago

Is there anyone who has recently passed the exam who can tell me approximately how many of the original questions are in the actual exam?

upvoted 4 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: A

A. Use S3 Event Notifications to write a message with image details to an Amazon Simple Queue Service (Amazon SQS) queue. Configure an AWS Lambda function to read the messages from the queue and to process the images.

upvoted 2 times

DanielWuTRT 11 months, 2 weeks ago

Selected Answer: A

We made it to the last one! Good luck!

upvoted 5 times

NSA_Poker 11 months, 3 weeks ago

Selected Answer: A

(B) is eliminated. Reserved instances are more expensive than Spot Fleet.
(C) is eliminated. Container instance more expensive than Lambda. SQS needed NOT SNS.
(D) is eliminated. Elastic Beanstalk is more expensive than Spot Fleet. SQS needed NOT SNS.
(A) is correct. It's the most cost effective service & the scope of its capabilities are within the requirements.

upvoted 5 times

trinh_le 1 year, 1 month ago

Selected Answer: A

less than 5 minutes -> use lambda

upvoted 6 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #905 *Topic 1*

A company wants to improve the availability and performance of its hybrid application. The application consists of a stateful TCP-based workload hosted on Amazon EC2 instances in different AWS Regions and a stateless UDP-based workload hosted on premises.

Which combination of actions should a solutions architect take to improve availability and performance? (Choose two.)

A. Create an accelerator using AWS Global Accelerator. Add the load balancers as endpoints. **Most Voted**

B. Create an Amazon CloudFront distribution with an origin that uses Amazon Route 53 latency-based routing to route requests to the load balancers.

C. Configure two Application Load Balancers in each Region. The first will route to the EC2 endpoints, and the second will route to the on-premises endpoints.

D. Configure a Network Load Balancer in each Region to address the EC2 endpoints. Configure a Network Load Balancer in each Region that routes to the on-premises endpoints. **Most Voted**

E. Configure a Network Load Balancer in each Region to address the EC2 endpoints. Configure an Application Load Balancer in each Region that routes to the on-premises endpoints.

Correct Answer: *AD*

Community vote distribution

AD (100%)

Comments

Ricksaurus Highly Voted 10 months, 1 week ago

Selected Answer: AD

TCP >> NLB
non-http >> accelerator.

upvoted 7 times

GOTJ Most Recent 3 months, 3 weeks ago

Selected Answer: AD

But, regarding to option "D", what's the reason behind the multi-regional deploy of the NLB for on-premises UDP endpoints? Shouldn't it be deployed on the nearest region only?

upvoted 2 times

muhammadahmer36 10 months ago

Selected Answer: AD

TCP >> NLB
non-http >> accelerator

upvoted 3 times

swati1508 10 months, 1 week ago

AD
-TCP use NLB
FOR non-http use accelerator.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #906 *Topic 1*

A company runs a self-managed Microsoft SQL Server on Amazon EC2 instances and Amazon Elastic Block Store (Amazon EBS). Daily snapshots are taken of the EBS volumes.

Recently, all the company's EBS snapshots were accidentally deleted while running a snapshot cleaning script that deletes all expired EBS snapshots. A solutions architect needs to update the architecture to prevent data loss without retaining EBS snapshots indefinitely.

Which solution will meet these requirements with the LEAST development effort?

A. Change the IAM policy of the user to deny EBS snapshot deletion.

B. Copy the EBS snapshots to another AWS Region after completing the snapshots daily.

C. Create a 7-day EBS snapshot retention rule in Recycle Bin and apply the rule for all snapshots. **Most Voted**

D. Copy EBS snapshots to Amazon S3 Standard-Infrequent Access (S3 Standard-IA).

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

FlyingHawk 4 months, 3 weeks ago

Selected Answer: C

AWS Recycle Bin allows you to set retention rules for EBS snapshots, preventing them from being permanently deleted for a specified period (e.g., 7 days).

If a snapshot is accidentally deleted, it is moved to the Recycle Bin and can be restored within the retention period.
https://docs.aws.amazon.com/ebs/latest/userguide/recycle-bin.html

upvoted 1 times

dhewa 9 months, 3 weeks ago

Selected Answer: C

Provides a safety net against accidental deletions

upvoted 1 times

progounick 9 months, 3 weeks ago

Selected Answer: C

C. Snapshot can be recovered if accidently deleted

upvoted 1 times

pujithacg8 10 months ago

C, AWS Recycle Bin allows you to recover resources like EBS snapshots that were accidentally deleted.

upvoted 1 times

AnasAWS 10 months, 1 week ago

Selected Answer: C

7 Days retention period will protect snapshots from being accedently perminantly deleted.

upvoted 1 times

Ricksaurus 10 months, 1 week ago

Selected Answer: C

Snapshot can be recovered if accidently deleted

upvoted 1 times

example_ 10 months, 1 week ago

Selected Answer: C

https://aws.amazon.com/blogs/aws/new-recycle-bin-for-ebs-snapshots/

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #907 *Topic 1*

A company wants to use an AWS CloudFormation stack for its application in a test environment. The company stores the CloudFormation template in an Amazon S3 bucket that blocks public access. The company wants to grant CloudFormation access to the template in the S3 bucket based on specific user requests to create the test environment. The solution must follow security best practices.

Which solution will meet these requirements?

A. Create a gateway VPC endpoint for Amazon S3. Configure the CloudFormation stack to use the S3 object URL.

B. Create an Amazon API Gateway REST API that has the S3 bucket as the target. Configure the CloudFormation stack to use the API Gateway URL.

C. Create a presigned URL for the template object. Configure the CloudFormation stack to use the presigned URL. **Most Voted**

D. Allow public access to the template object in the S3 bucket. Block the public access after the test environment is created.

Correct Answer: *C*

Community vote distribution

C (71%) A (29%)

Comments

AnasAWS Highly Voted 10 months, 1 week ago

Selected Answer: C

https://docs.aws.amazon.com/AmazonS3/latest/userguide/ShareObjectPreSignedURL.html

upvoted 6 times

blehbleh 9 months ago

Thank you for posting this.

upvoted 2 times

Dantecito Most Recent 3 months, 2 weeks ago

Selected Answer: C

C. Not the most secure but the only that works.
A. Cloudformation is a regional service so is outside the vpc and gateway VPC endpoint only works inside a VPC.
B. Doesn't do anything, we need a lambda between api gateway and cloudformation.
D. public access is a no.

upvoted 2 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: A

OptionA the gateway endpoint for S3 provides private access to the S3 template, aligns with security best practices, and incurs no additional costs. Option C (presigned URLs) potentially exposing the template to the public internet, even if temporarily. You only use C for the following reasons If you need a quick, temporary solution for granting access to the template.

If you are not using a VPC or do not want to set up a VPC endpoint.

If the template is not highly sensitive and temporary public exposure is acceptable.

upvoted 2 times

FlyingHawk 4 months, 3 weeks ago

Given the requirement to follow security best practices, Option A (VPC Endpoint) is the better choice. While it requires slightly more effort to set up initially, it provides a secure, private, and persistent solution that aligns with security best practices.

If development effort is a critical factor and the company is comfortable with the temporary public exposure of the template, Option C (Presigned URL) can be considered as a quick and easy alternative.

upvoted 1 times

robotgeek 5 months, 3 weeks ago

Selected Answer: A

The question states "specific user requests" meaning a shared URL does not meet the requirement. BUT with gateway VPC endpoint you can specify custom policies as in the following link
https://docs.aws.amazon.com/vpc/latest/privatelink/vpc-endpoints-access.html

upvoted 2 times

pujithacg8 10 months ago

C, A presigned URL grants temporary access to an S3 object without making it publicly accessible.

upvoted 1 times

flaviobrf 10 months, 1 week ago

Selected Answer: C

For me C is the right answer

upvoted 2 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #908 *Topic 1*

A company has applications that run in an organization in AWS Organizations. The company outsources operational support of the applications. The company needs to provide access for the external support engineers without compromising security.

The external support engineers need access to the AWS Management Console. The external support engineers also need operating system access to the company's fleet ofAmazon EC2 instances that run Amazon Linux in private subnets.

Which solution will meet these requirements MOST securely?

A. Confirm that AWS Systems Manager Agent (SSM Agent) is installed on all instances. Assign an instance profile with the necessary policy to connect to Systems Manager. Use AWS IAM Identity Center to provide the external support engineers console access. Use Systems Manager Session Manager to assign the required permissions. **Most Voted**

B. Confirm that AWS Systems Manager Agent (SSM Agent) is installed on all instances. Assign an instance profile with the necessary policy to connect to Systems Manager. Use Systems Manager Session Manager to provide local IAM user credentials in each AWS account to the external support engineers for console access.

C. Confirm that all instances have a security group that allows SSH access only from the external support engineers' source IP address ranges. Provide local IAM user credentials in each AWS account to the external support engineers for console access. Provide each external support engineer an SSH key pair to log in to the application instances.

D. Create a bastion host in a public subnet. Set up the bastion host security group to allow access from only the external engineers' IP address ranges. Ensure that all instances have a security group that allows SSH access from the bastion host. Provide each external support engineer an SSH key pair to log in to the application instances. Provide local account IAM user credentials to the engineers for console access.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

LeonSauveterre 5 months ago

Selected Answer: A

A - It leverages AWS-managed solutions that simplify operations, enhance security, and meet the requirement for both console and OS-level access, while avoiding the pitfalls of SSH-based access or local user management, so this is what we're looking for

and OS-level access, while avoiding the pitfalls of SSH-based access or local user management, so this is what we're looking for.
B - Providing local IAM user credentials increases operational overhead and security risks. Managing these credentials across multiple AWS accounts can lead to inconsistency and potential vulnerabilities.
C - Allowing SSH access via security groups and distributing SSH key pairs introduces significant security risks. Keys could be lost or misused, and managing IP address ranges is not a piece of cake.
D - Bastion hosts require ongoing maintenance and monitoring to ensure security, which is unnecessary with Session Manager.

upvoted 2 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: A

B: I believe the Session Manager may not have the capability to provide IAM user credentials. If it were to do so, it could potentially lead to security concerns.

upvoted 1 times

Cpso 6 months, 1 week ago

Selected Answer: A

A or B. Exam not tell if they have IDP and no. of outsource.

upvoted 2 times

Changwha 6 months, 2 weeks ago

A: This solution minimizes the attack surface by eliminating the need for SSH access, avoids bastion hosts, and provides secure, auditable, and federated access to AWS resources.

upvoted 2 times

swati1508 10 months ago

A use MSM

upvoted 2 times

officedepotadmin 10 months, 1 week ago

Selected Answer: A

Systems Manager Session Manager allows secure, auditable, and controlled access to your EC2 instances without needing to open SSH ports or manage SSH keys, reducing the attack surface.
Local IAM user credentials are less secure and harder to manage at scale compared to using IAM Identity Center.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #909 *Topic 1*

A company uses Amazon RDS for PostgreSQL to run its applications in the us-east-1 Region. The company also uses machine learning (ML) models to forecast annual revenue based on near real-time reports. The reports are generated by using the same RDS for PostgreSQL database. The database performance slows during business hours. The company needs to improve database performance.

Which solution will meet these requirements MOST cost-effectively?

A. Create a cross-Region read replica. Configure the reports to be generated from the read replica.

B. Activate Multi-AZ DB instance deployment for RDS for PostgreSQL. Configure the reports to be generated from the standby database.

C. Use AWS Data Migration Service (AWS DMS) to logically replicate data to a new database. Configure the reports to be generated from the new database.

D. Create a read replica in us-east-1. Configure the reports to be generated from the read replica. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

hpirnaj 5 months ago

Selected Answer: D

MOST cost-effectively

upvoted 1 times

[Removed] 9 months, 3 weeks ago

Selected Answer: D

D. Create a read replica in us-east-1. Configure the reports to be generated from the read replica.

upvoted 1 times

muhammadahmer36 10 months ago

Selected Answer: D

Read replicas are typically less expensive than setting up a cross-Region replica or activating Multi-AZ deployments. You only pay for the additional read replica, without the overhead costs associated with cross-Region data transfer or maintaining a synchronous standby in Multi-AZ setups.

upvoted 1 times

officedepotadmin 10 months, 1 week ago

Selected Answer: D

Read replicas are typically less expensive than setting up a cross-Region replica or activating Multi-AZ deployments. You only pay for the additional read replica, without the overhead costs associated with cross-Region data transfer or maintaining a synchronous standby in Multi-AZ setups.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #910 *Topic 1*

A company hosts its multi-tier, public web application in the AWS Cloud. The web application runs on Amazon EC2 instances, and its database runs on Amazon RDS. The company is anticipating a large increase in sales during an upcoming holiday weekend. A solutions architect needs to build a solution to analyze the performance of the web application with a granularity of no more than 2 minutes.

What should the solutions architect do to meet this requirement?

A. Send Amazon CloudWatch logs to Amazon Redshift. Use Amazon QuickS ght to perform further analysis.

B. Enable detailed monitoring on all EC2 instances. Use Amazon CloudWatch metrics to perform further analysis. Most Voted

C. Create an AWS Lambda function to fetch EC2 logs from Amazon CloudWatch Logs. Use Amazon CloudWatch metrics to perform further analysis.

D. Send EC2 logs to Amazon S3. Use Amazon Redshift to fetch logs from the S3 bucket to process raw data for further analysis with Amazon QuickSight.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

[Removed] 9 months, 3 weeks ago

Selected Answer: B

B. Enable detailed monitoring on all EC2 instances. Use Amazon CloudWatch metrics to perform further analysis.

upvoted 2 times

muhammadahmer36 10 months ago

Selected Answer: B

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-cloudwatch-new.html

upvoted 2 times

officedepotadmin 10 months, 1 week ago

Selected Answer: B

Enabling detailed monitoring on EC2 instances provides metrics at a 1-minute granularity, which is well within the required 2-minute granularity for performance analysis.

upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: B

Answer is B

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #911 *Topic 1*

A company runs an application that stores and shares photos. Users upload the photos to an Amazon S3 bucket. Every day, users upload approximately 150 photos. The company wants to design a solution that creates a thumbnail of each new photo and stores the thumbnail in a second S3 bucket.

Which solution will meet these requirements MOST cost-effectively?

A. Configure an Amazon EventBridge scheduled rule to invoke a script every minute on a long-running Amazon EMR cluster. Configure the script to generate thumbnails for the photos that do not have thumbnails. Configure the script to upload the thumbnails to the second S3 bucket.

B. Configure an Amazon EventBridge scheduled rule to invoke a script every minute on a memory-optimized Amazon EC2 instance that is always on. Configure the script to generate thumbnails for the photos that do not have thumbnails. Configure the script to upload the thumbnails to the second S3 bucket.

C. Configure an S3 event notification to invoke an AWS Lambda function each time a user uploads a new photo to the application. Configure the Lambda function to generate a thumbnail and to upload the thumbnail to the second S3 bucket. **Most Voted**

D. Configure S3 Storage Lens to invoke an AWS Lambda function each time a user uploads a new photo to the application. Configure the Lambda function to generate a thumbnail and to upload the thumbnail to a second S3 bucket.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

ieffiong 3 months, 1 week ago

Selected Answer: C

Take a sip of water and smile, you're almost there.

upvoted 1 times

[Removed] 9 months, 3 weeks ago

Selected Answer: C

C. Configure an S3 event notification to invoke an AWS Lambda function each time a user uploads a new photo to the

application. Configure the Lambda function to generate a thumbnail and to upload the thumbnail to the second S3 bucket.
upvoted 3 times

swati1508 10 months ago

C is correct
upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: C

Answer is C
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #912 *Topic 1*

A company has stored millions of objects across multiple prefixes in an Amazon S3 bucket by using the Amazon S3 Glacier Deep Archive storage class. The company needs to delete all data older than 3 years except for a subset of data that must be retained. The company has identified the data that must be retained and wants to implement a serverless solution.

Which solution will meet these requirements?

A. Use S3 Inventory to list all objects. Use the AWS CLI to create a script that runs on an Amazon EC2 instance that deletes objects from the inventory list.

B. Use AWS Batch to delete objects older than 3 years except for the data that must be retained.

C. Provision an AWS Glue crawler to query objects older than 3 years. Save the manifest file of old objects. Create a script to delete objects in the manifest.

D. Enable S3 Inventory. Create an AWS Lambda function to filter and delete objects. Invoke the Lambda function with S3 Batch Operations to delete objects by using the inventory reports. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

Salilgen 5 months ago

Selected Answer: D

https://docs.aws.amazon.com/AmazonS3/latest/userguide/batch-ops.html

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: D

Takeaways:
1. S3 Inventory provides a flat file report of the objects in an S3 bucket, including metadata such as storage class, size, and last modified date. Now you can easily find files "older than 3 years".
2. The Lambda function can be used to filter objects based on the company's requirements (excluding the subset of data to retain) and to delete the objects programmatically.
3. S3 Batch Operations allow the execution of bulk operations (e.g., delete operations) on a large number of S3 objects.

A - Running a script on an EC2 instance is not serverless.
B - AWS Batch is designed for batch computing workloads, not operations like deleting S3 objects.
C - AWS Glue crawlers are designed to extract metadata and schema from data stored in S3 for querying and ETL processes, not for managing or deleting objects.

upvoted 2 times

agbor_tambe 8 months, 1 week ago

yes D is the right answer

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: D

D sounds right

upvoted 2 times

muhammadahmer36 10 months ago

Selected Answer: D

Enable S3 Inventory. Create an AWS Lambda function to filter and delete objects. Invoke the Lambda function with S3 Batch Operations to delete objects by using the inventory reports.

upvoted 3 times

swati1508 10 months ago

D is correct

upvoted 2 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #913 *Topic 1*

A company is building an application on AWS. The application uses multiple AWS Lambda functions to retrieve sensitive data from a single Amazon S3 bucket for processing. The company must ensure that only authorized Lambda functions can access the data. The solution must comply with the principle of least privilege.

Which solution will meet these requirements?

A. Grant full S3 bucket access to all Lambda functions through a shared IAM role.

B. Configure the Lambda functions to run within a VPC. Configure a bucket policy to grant access based on the Lambda functions' VPC endpoint IP addresses.

C. Create individual IAM roles for each Lambda function. Grant the IAM roles access to the S3 bucket. Assign each IAM role as the Lambda execution role for its corresponding Lambda function. **Most Voted**

D. Configure a bucket policy granting access to the Lambda functions based on their function ARNs.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

LeonSauveterre 5 months ago

Selected Answer: C

A - Full access? No!
B - Theoretically possible but IP-based policies are harder to manage and not as secure as IAM roles.
C - By creating separate IAM roles for each Lambda function, permissions can be narrowly scoped to each function's specific needs.
D - This would actually work. It's just that bucket policies based on function ARNs are less flexible and scalable compared to IAM roles, especially when you need to manage permissions for multiple Lambda functions or add new functions.
upvoted 1 times

LeonSauveterre 5 months ago

Out of my instinct, least privilege leads to individually managing each function.
upvoted 1 times

56ce46c 8 months, 3 weeks ago

i think D is also right
S3 Bucket Policy: Use an S3 bucket policy that grants access to the specific Lambda functions based on their function ARNs. This ensures that only the authorized Lambda functions can retrieve data from the S3 bucket.

upvoted 3 times

[Removed] 9 months, 3 weeks ago

Selected Answer: C

C sounds right

upvoted 2 times

swati1508 10 months ago

A, B and D wrong only C is right

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #914 *Topic 1*

A company has developed a non-production application that is composed of multiple microservices for each of the company's business units. A single development team maintains all the microservices.

The current architecture uses a static web frontend and a Java-based backend that contains the application logic. The architecture also uses a MySQL database that the company hosts on an Amazon EC2 instance.

The company needs to ensure that the application is secure and available globally.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use Amazon CloudFront and AWS Amplify to host the static web frontend. Refactor the microservices to use AWS Lambda functions that the microservices access by using Amazon API Gateway. Migrate the MySQL database to an Amazon EC2 Reserved Instance.

B. Use Amazon CloudFront and Amazon S3 to host the static web frontend. Refactor the microservices to use AWS Lambda functions that the microservices access by using Amazon API Gateway. Migrate the MySQL database to Amazon RDS for MySQL. **Most Voted**

C. Use Amazon CloudFront and Amazon S3 to host the static web frontend. Refactor the microservices to use AWS Lambda functions that are in a target group behind a Network Load Balancer. Migrate the MySQL database to Amazon RDS for MySQL.

D. Use Amazon S3 to host the static web frontend. Refactor the microservices to use AWS Lambda functions that are in a target group behind an Application Load Balancer. Migrate the MySQL database to an Amazon EC2 Reserved Instance.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

[Removed] 9 months, 3 weeks ago

Selected Answer: B

B sounds right

upvoted 2 times

upvoted 2 times

muhammadahmer36 10 months ago

Selected Answer: B

Answer is B

upvoted 2 times

komorebi 10 months ago

Selected Answer: B

Answer is B

upvoted 3 times

swati1508 10 months ago

Answer is D

upvoted 1 times

swati1508 10 months ago

B is correct sorry

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #915 *Topic 1*

A video game company is deploying a new gaming application to its global users. The company requires a solution that will provide near real-time reviews and rankings of the players.

A solutions architect must design a solution to provide fast access to the data. The solution must also ensure the data persists on disks in the event that the company restarts the application.

Which solution will meet these requirements with the LEAST operational overhead?

A. Configure an Amazon CloudFront distribution with an Amazon S3 bucket as the origin. Store the player data in the S3 bucket.

B. Create Amazon EC2 instances in multiple AWS Regions. Store the player data on the EC2 instances. Configure Amazon Route 53 with geolocation records to direct users to the closest EC2 instance.

C. Deploy an Amazon ElastiCache for Redis duster. Store the player data in the ElastiCache cluster. **Most Voted**

D. Deploy an Amazon ElastiCache for Memcached duster. Store the player data in the ElastiCache cluster.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

officedepotadmin Highly Voted 10 months ago

Selected Answer: C

Amazon ElastiCache for Redis provides in-memory caching which ensures low latency and high throughput, perfect for near real-time access to player reviews and rankings.

Redis supports data persistence by snapshotting data to disk (RDB snapshots) and appending changes to a log (AOF), ensuring that the data is not lost even if the application restarts.

upvoted 5 times

Salilgen Most Recent 5 months ago

Selected Answer: C

https://aws.amazon.com/blogs/database/building-a-real-time-gaming-leaderboard-with-amazon-elasticache-for-redis/

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: C

To be more exact, you should use Amazon ElastiCache for Redis with AOF (Append-Only File) persistence. BTW, all the "duster"s in the options should be "cluster"s.

upvoted 1 times

Bwhizzy 7 months, 4 weeks ago

Selected Answer: C

REDIS
• Multi AZ with Auto-Failover
• Read Replicas to scale reads and have high availability
• Data Durability using AOF persistence
• Backup and restore features
• Supports Sets and Sorted Sets

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: C

C sounds right

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #916 *Topic 1*

A company is designing an application on AWS that processes sensitive data. The application stores and processes financial data for multiple customers.

To meet compliance requirements, the data for each customer must be encrypted separately at rest by using a secure, centralized key management solution. The company wants to use AWS Key Management Service (AWS KMS) to implement encryption.

Which solution will meet these requirements with the LEAST operational overhead?

A. Generate a unique encryption key for each customer. Store the keys in an Amazon S3 bucket. Enable server-side encryption.

B. Deploy a hardware security appliance in the AWS environment that securely stores customer-provided encryption keys. Integrate the security appliance with AWS KMS to encrypt the sensitive data in the application.

C. Create a single AWS KMS key to encrypt all sensitive data across the application.

D. Create separate AWS KMS keys for each customer's data that have granular access control and logging enabled.
Most Voted

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

officedepotadmin Highly Voted 10 months ago

Selected Answer: D

While enabling server-side encryption in S3 can manage encryption, it does not offer the same level of control and auditing as AWS KMS. Managing individual keys manually in S3 would also increase operational overhead.

upvoted 7 times

Jeyaluxshan Most Recent 9 months ago

D is with less management overhead

upvoted 2 times

dhewa 9 months, 3 weeks ago

Selected Answer: D

D is more secure

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: D

D sounds right

upvoted 2 times

progounick 9 months, 3 weeks ago

Selected Answer: D

it is obvious that D is correct

upvoted 1 times

muhammadahmer36 10 months ago

Selected Answer: D

hile enabling server-side encryption in S3 can manage encryption, it does not offer the same level of control and auditing as AWS KMS. Managing individual keys manually in S3 would also increase operational overhead.

upvoted 2 times

nebajp 10 months ago

Selected Answer: D

D is the correct Answer

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #917 *Topic 1*

A company needs to design a resilient web application to process customer orders. The web application must automatically handle increases in web traffic and application usage without affecting the customer experience or losing customer orders.

Which solution will meet these requirements?

A. Use a NAT gateway to manage web traffic. Use Amazon EC2 Auto Scaling groups to receive, process, and store processed customer orders. Use an AWS Lambda function to capture and store unprocessed orders.

B. Use a Network Load Balancer (NLB) to manage web traffic. Use an Application Load Balancer to receive customer orders from the NLUse Amazon Redshift with a Multi-AZ deployment to store unprocessed and processed customer orders.

C. Use a Gateway Load Balancer (GWLB) to manage web traffic. Use Amazon Elastic Container Service (Amazon ECS) to receive and process customer orders. Use the GWLB to capture and store unprocessed orders. Use Amazon DynamoDB to store processed customer orders.

D. Use an Application Load Balancer to manage web traffic. Use Amazon EC2 Auto Scaling groups to receive and process customer orders. Use Amazon Simple Queue Service (Amazon SQS) to store unprocessed orders. Use Amazon RDS with a Multi-AZ deployment to store processed customer orders. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

Cpso 6 months, 1 week ago

Selected Answer: D

D is right. But explaining is unclear. should have 2 group of ec2. one for accept order to queue. another to process and store to DB.

upvoted 1 times

[Removed] 9 months, 3 weeks ago

Selected Answer: D

D sounds right

upvoted 2 times

pujithacg8 10 months ago

D is perfect

upvoted 2 times

komorebi 10 months ago

Selected Answer: D

Answer is D

upvoted 3 times

swati1508 10 months ago

Answer is D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #918 *Topic 1*

A company is using AWS DataSync to migrate millions of files from an on-premises system to AWS. The files are 10 KB in size on average.

The company wants to use Amazon S3 for file storage. For the first year after the migration, the files will be accessed once or twice and must be immediately available. After 1 year, the files must be archived for at least 7 years.

Which solution will meet these requirements MOST cost-effectively?

A. Use an archive tool to group the files into large objects. Use DataSync to migrate the objects. Store the objects in S3 Glacier Instant Retrieval for the first year. Use a lifecycle configuration to transition the files to S3 Glacier Deep Archive after 1 year with a retention period of 7 years. **Most Voted**

B. Use an archive tool to group the files into large objects. Use DataSync to copy the objects to S3 Standard-Infrequent Access (S3 Standard-IA). Use a lifecycle configuration to transition the files to S3 Glacier Instant Retrieval after 1 year with a retention period of 7 years.

C. Configure the destination storage class for the files as S3 Glacier Instant Retrieval. Use a lifecycle policy to transition the files to S3 Glacier Flexible Retrieval after 1 year with a retention period of 7 years.

D. Configure a DataSync task to transfer the files to S3 Standard-Infrequent Access (S3 Standard-IA). Use a lifecycle configuration to transition the files to S3 Deep Archive after 1 year with a retention period of 7 years.

Correct Answer: *A*

Community vote distribution

A (54%) D (45%)

Comments

dhewa Highly Voted 9 months, 3 weeks ago

Selected Answer: D

D simplifies the process by directly using S3 Standard-IA and then transitioning to S3 Glacier Deep Archive, which aligns well with access patterns and cost requirements. Option A sounds right but using an archive tool to group files into large objects adds complexity and operational overhead. This step isn't necessary if you can directly manage the files with S3 lifecycle policies.

upvoted 12 times

GOTJ 3 months, 3 weeks ago

Correct during the first year (online period). Incorrect during years 2-8 (archive period). Go check the numbers in price calculator yourself with the following scenarios for the archive period (S3 Glacier Deep Archive):

* For small files: 100,000,000 files of 10KB (1000GB aprox). Considering both, the transition cost of 100,000,000 files to the Glacier Deep Archive class and the maintenance cost of 7 years for those 100,000,000 files: $6,940,34
* For big files: 1000 files of 1000MB (1000GB aprox). Considering both, the transition cost of 1,000 files to the Glacier Deep Archive class and the maintenance cost of 7 years for those 1,000 files: $83,23

upvoted 2 times

blehbleh Highly Voted 9 months ago

Selected Answer: A

Its A, took some research but A is correct Per Amazon S# glacier page: "Amazon S3 Glacier Instant Retrieval is an archive storage class that delivers the lowest-cost storage for long-lived data that is rarely accessed and requires retrieval in milliseconds. With S3 Glacier Instant Retrieval, you can save up to 68% on storage costs compared to using the S3 Standard-Infrequent Access (S3 Standard-IA) storage class, when your data is accessed once per quarter." After the one year move it to Deep archive.

upvoted 6 times

TEC65 Most Recent 1 month, 1 week ago

Selected Answer: D

def d is the answer

upvoted 1 times

ad2dj28 2 months, 3 weeks ago

Selected Answer: D

10/10 TIMES STANDARD INFREQUENT ACCESS IS CHEAPER THAN USING GLACIER INSTANT RETRIEVAL WHEN COMPARING COSTS WITH PRICING CALCULATOR

upvoted 1 times

tch 3 months ago

Selected Answer: A

B & C are out

upvoted 1 times

tch 3 months ago

Amazon S3 Glacier Instant Retrieval is an archive storage class that delivers the lowest-cost storage for long-lived data that is rarely accessed and requires retrieval in milliseconds.

With Amazon S3 Glacier Instant Retrieval, you can save up to 68% on storage costs compared to using the S3 Standard-Infrequent Access (S3 Standard-IA) storage class, when your data is accessed once per quarter." After the one year move it to S3 Glacier Deep archive.

upvoted 1 times

Mistwalker 4 months, 2 weeks ago

Selected Answer: A

10KB/file x millions of files = insane unnecessary cost when each is charged at 128KB/object
No one is talking about operational overhead, they just want cost-effectiveness

upvoted 2 times

GOTJ 3 months, 3 weeks ago

Correct, but they mention that the files must be "immediately available" during the first year. So any option involving archive tools should be discarded for this period, due to the extra steps needed to accessing the files: figuring out which archive contains the file you are looking for, download the archive file and extract the desired file. I think this simple requirement explains the almost 50%-50% consensus for this particular question.

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: A

- Each file is 10 KB in size.
- S3 charges a minimum of 128 KB per object.
-Total number of files: millions of files.
- Files are accessed twice in the first year.
- After 1 year, files are archived for 7 years
- S3 Glacier: Minimum charge for 40 KB per object.
S3 Standard-IA and S3 One Zone-IA storage have a minimum billable object size of 128 KB. Smaller objects may be stored but will be charged for 128 KB of storage at the appropriate storage class rate Option D does not group those files into big object, so 1 million files will be charged as 1 million objects(128KB).
https://aws.amazon.com/s3/pricing/
upvoted 3 times

FlyingHawk 4 months, 3 weeks ago

Note
The S3 Standard-IA and S3 One Zone-IA storage classes are suitable for objects larger than 128 KB that you plan to store for at least 30 days. If an object is less than 128 KB, Amazon S3 charges you for 128 KB. If you delete an object before the end of the 30-day minimum storage duration period, you are charged for 30 days. Objects that are deleted, overwritten, or transitioned to a different storage class before 30 days will incur the normal storage usage charge plus a pro-rated charge for the remainder of the 30-day minimum. For pricing information, see Amazon S3 pricing.
https://docs.aws.amazon.com/AmazonS3/latest/userguide/storage-class-intro.html
upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: D

About option A: For data stored in the S3 Glacier Instant Retrieval storage class, there is indeed a minimum object size of 128 KB. If the objects are smaller than 128 KB, AWS still charges for 128 KB per object, which can significantly impact cost efficiency when dealing with a large number of small files, such as the 10 KB files in the given scenario.

One more thing: once or twice a year is perfect for Infrequent Access. It also provides instant availability for the objects stored in it.
upvoted 1 times

LeonSauveterre 5 months ago

I believe Standard-IA remains the most straightforward and cost-effective choice for small files (10 KB average) that are accessed 1~2 times in the first year without aggregation.
upvoted 1 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: A

S3 Glacier Instant Retrieval is more expensive than S3 Standard-IA for the first year when immediate availability is required.
upvoted 2 times

EllenLiu 5 months, 2 weeks ago

sorry, choose D
upvoted 1 times

dragossky 5 months, 3 weeks ago

Selected Answer: A

For sure A, has amazing retrieval - Store the objects in S3 Glacier Instant Retrieval
upvoted 2 times

jfedotov 5 months, 3 weeks ago

Selected Answer: D

D is correct.
A is not correct, because Amazon S3 Glacier Instant Retrieval min object size is 128KB, in question 10KB
upvoted 2 times

AMEJack 6 months, 1 week ago

Selected Answer: D

There is no need to use archive tool to migrate small files, DataSync is enough.

upvoted 2 times

Ez1tap 8 months, 2 weeks ago

Selected Answer: A

correct answer is A you need 2 separate lifecycle policy

upvoted 4 times

Abhiiinav 8 months, 4 weeks ago

Selected Answer: D

answer is D. Storing the objects in S3 Glacier Instant Retrieval for the first year is more expensive than S3 Standard-IA for data that is accessed infrequently.

upvoted 3 times

blehbleh 8 months ago

Wrong, it says 1-2 times in the first year and Amazon states that s3 glacier isn't at retrieval can save up to 68% compared to s3 infrequent access.

upvoted 2 times

Abdullah2004 9 months, 1 week ago

Selected Answer: D

D is most cost effective

upvoted 3 times

progounick 9 months, 1 week ago

Selected Answer: A

ChatGPT agrees with me and selected A

upvoted 2 times

youkarthik 6 months ago

My intance of chatgpt says it is D

upvoted 3 times

dhewa 9 months, 3 weeks ago

Selected Answer: A

D simplifies the process by directly using S3 Standard-IA and then transitioning to S3 Glacier Deep Archive, which aligns well with access patterns and cost requirements. Option A sounds right but using an archive tool to group files into large objects adds complexity and operational overhead. This step isn't necessary if you can directly manage the files with S3 lifecycle policies.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #919 *Topic 1*

A company recently performed a lift and shift migration of its on-premises Oracle database workload to run on an Amazon EC2 memory optimized Linux instance. The EC2 Linux instance uses a 1 TB Provisioned IOPS SSD (io1) EBS volume with 64,000 IOPS.

The database storage performance after the migration is slower than the performance of the on-premises database.

Which solution will improve storage performance?

A. Add more Provisioned IOPS SSD (io1) EBS volumes. Use OS commands to create a Logical Volume Management (LVM) stripe. **Most Voted**

B. Increase the Provisioned IOPS SSD (io1) EBS volume to more than 64,000 IOPS.

C. Increase the size of the Provisioned IOPS SSD (io1) EBS volume to 2 TB.

D. Change the EC2 Linux instance to a storage optimized instance type. Do not change the Provisioned IOPS SSD (io1) EBS volume.

Correct Answer: *A*

Community vote distribution

A (91%) B (9%)

Comments

FlyingHawk 4 months, 3 weeks ago

Selected Answer: A

https://docs.aws.amazon.com/prescriptive-guidance/latest/sql-server-ec2-best-practices/striping.html

upvoted 2 times

FlyingHawk 4 months, 3 weeks ago

When the performance requirement cannot be met by a single EBS volume, consider striping using Logical Volume Management (LVM). For example, if a single volume has 250 MiB/s throughput capacity, having a stripe set across four volumes can deliver 1000 MiB/s throughput.

Volumes should be of the same size and performance characteristics.
https://docs.aws.amazon.com/wellarchitected/latest/sap-lens/best-practice-14-2.html

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: A

A - For example, given two io1 volumes, each provisioned with 64,000 IOPS, would provide a combined 128,000 IOPS.
B - The maximum IOPS for an io1 volume is capped at 64,000 IOPS when attached to an EC2 instance.
C - While increasing the size of the volume increases baseline throughput (3 IOPS per GB), it won't exceed the 64,000 IOPS limit.
D - Storage optimized instances (like i3 or i4i) use instance store volumes, which are ephemeral (non-persistent) and cannot be used as a replacement for durable EBS volumes in this use case.

What is the deal with LVM striping?
LVM striping allows multiple volumes to be treated as a single logical volume. Reads and writes are distributed across the striped volumes, effectively increasing the available IOPS and throughput beyond the single-volume limits.

upvoted 4 times

JA2018 5 months, 2 weeks ago

Selected Answer: B

why I would choose Option B:

- Issue:
The problem is that the current IOPS (Input/Output Operations Per Second) provisioned on the EBS volume are not sufficient to meet the database workload demands, leading to slower performance compared to the on-premises setup.

- Solution:
Since the bottleneck is clearly identified as inadequate IOPS, the most effective solution is to simply increase the provisioned IOPS on the existing EBS volume, allowing for faster data read/write operations.

upvoted 1 times

jingen11 7 months, 3 weeks ago

Selected Answer: A

A

upvoted 2 times

elmyth 9 months ago

Selected Answer: A

Increase the Provisioned IOPS SSD (io1) EBS volume to more than 64,000 IOPS.

upvoted 2 times

pujithacg8 10 months ago

A is correct, The maximum provisioned IOPS for io1 is 64000 and hence you can achieve higher aggregate performance by adding more io1 volumes

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #920 *Topic 1*

A company is migrating from a monolithic architecture for a web application that is hosted on Amazon EC2 to a serverless microservices architecture. The company wants to use AWS services that support an event-driven, loosely coupled architecture. The company wants to use the publish/subscribe (pub/sub) pattern.

Which solution will meet these requirements MOST cost-effectively?

A. Configure an Amazon API Gateway REST API to invoke an AWS Lambda function that publishes events to an Amazon Simple Queue Service (Amazon SQS) queue. Configure one or more subscribers to read events from the SQS queue.

B. Configure an Amazon API Gateway REST API to invoke an AWS Lambda function that publishes events to an Amazon Simple Notification Service (Amazon SNS) topic. Configure one or more subscribers to receive events from the SNS topic.

C. Configure an Amazon API Gateway WebSocket API to write to a data stream in Amazon Kinesis Data Streams with enhanced fan-out. Configure one or more subscribers to receive events from the data stream.

D. Configure an Amazon API Gateway HTTP API to invoke an AWS Lambda function that publishes events to an Amazon Simple Notification Service (Amazon SNS) topic. Configure one or more subscribers to receive events from the topic.
Most Voted

Correct Answer: *D*

Community vote distribution

D (74%) B (26%)

Comments

srcntpc 3 weeks ago
Selected Answer: D

D for sure
upvoted 1 times

bora4motion 3 weeks, 3 days ago
Selected Answer: D

I hate so much the "most cost effectively" questions. While B is also valid from a technical perspective apparently D is cheaper.
upvoted 2 times

tch 3 months, 1 week ago

REST APIs and HTTP APIs are both RESTful API products. REST APIs support more features than HTTP APIs, while HTTP APIs are designed with minimal features so that they can be offered at a lower price. Choose REST APIs if you need features such as API keys, per-client throttling, request validation, AWS WAF integration, or private API endpoints. Choose HTTP APIs if you don't need the features included with REST APIs.

upvoted 2 times

tomasmolinari 4 months, 2 weeks ago

Selected Answer: D

"HTTP APIs enable you to create RESTful APIs with lower latency and lower cost than REST APIs."
HTTP APIs are more cost-effective than REST APIs

upvoted 3 times

LeonSauveterre 5 months ago

Selected Answer: B

Amazon SNS is perfectly designed to support the pub/sub pattern efficiently. So B or D.
About option D: Using an HTTP API instead of a REST API is less common when using SNS for event publishing. Typically, REST APIs in API Gateway are used for interacting with Lambda functions, making option B a more conventional choice.

That said, scenarios vary based on your actual, detailed needs, like predicted frequencies and costs. For the SAA exam, I'll definitely go with B.

upvoted 1 times

jfedotov 5 months ago

Selected Answer: D

Http api

upvoted 1 times

Denise123 5 months, 1 week ago

Selected Answer: B

While the Amazon API Gateway HTTP API (Option D) can be more cost-effective for simple APIs with low data transfer and minimal feature requirements, the Amazon API Gateway REST API (Option B) may still be more cost-effective in many real-world scenarios, especially for more complex web applications with additional requirements.

The choice between Options B and D should be based on a careful evaluation of your specific requirements, existing infrastructure, and potential cost savings or trade-offs involved.

upvoted 1 times

78b9037 5 months, 3 weeks ago

Selected Answer: D

HTTP API is cheaper. The difference in pricing is steep. REST APIs will run you USD $3.50 per one million requests plus charges for data transferred out. By contrast, HTTP APIs cost a mere $1.00 per request for the first million requests and then $0.90 per million requests after that.

upvoted 3 times

jingen11 7 months, 3 weeks ago

Selected Answer: D

Question ask for cheapest.
Http api offers everything the requirements asked for. i.e Lambda function for serverless, event driven.

upvoted 2 times

fis_examtopic 8 months, 4 weeks ago

HTTP APIs with Lambda. I'm go with D

upvoted 1 times

progounick 9 months, 2 weeks ago

Selected Answer: B

rest is cheaper than http

upvoted 1 times

Noveo 9 months ago

Quite the opposite: "REST APIs support more features than HTTP APIs, while HTTP APIs are designed with minimal features so that they can be offered at a lower price."

upvoted 3 times

Denise123 5 months, 1 week ago

The Amazon API Gateway HTTP API is a more recent addition to API Gateway and is designed to be a more cost-effective option for simple APIs. However, it has fewer features and capabilities compared to the REST API.
Here's a comparison of the costs between the two API Gateway options:
-REST API: $3.50 per million API calls received, plus data transfer costs.
-HTTP API: $0.90 per million API calls received, plus data transfer costs.
As you can see, the HTTP API is cheaper in terms of the API call pricing. However, the REST API may still be more cost-effective in certain scenarios. I am not sure what is needed in this question though.

upvoted 2 times

mk168898 6 months, 4 weeks ago

based on this answer should be D

upvoted 1 times

dhewa 9 months, 3 weeks ago

Selected Answer: B

An HTTP API instead might not be necessary for this use case.

upvoted 2 times

muhammadahmer36 10 months ago

Selected Answer: D

D is the right answer

upvoted 2 times

pujithacg8 10 months ago

B or D, Will go with B

upvoted 2 times

muhammadahmer36 10 months ago

Why B, not D?

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #921 *Topic 1*

A company recently migrated a monolithic application to an Amazon EC2 instance and Amazon RDS. The application has tightly coupled modules. The existing design of the application gives the application the ability to run on only a single EC2 instance.

The company has noticed high CPU utilization on the EC2 instance during peak usage times. The high CPU utilization corresponds to degraded performance on Amazon RDS for read requests. The company wants to reduce the high CPU utilization and improve read request performance.

Which solution will meet these requirements?

A. Resize the EC2 instance to an EC2 instance type that has more CPU capacity. Configure an Auto Scaling group with a minimum and maximum size of 1. Configure an RDS read replica for read requests. **Most Voted**

B. Resize the EC2 instance to an EC2 instance type that has more CPU capacity. Configure an Auto Scaling group with a minimum and maximum size of 1. Add an RDS read replica and redirect all read/write traffic to the replica.

C. Configure an Auto Scaling group with a minimum size of 1 and maximum size of 2. Resize the RDS DB instance to an instance type that has more CPU capacity.

D. Resize the EC2 instance to an EC2 instance type that has more CPU capacity. Configure an Auto Scaling group with a minimum and maximum size of 1. Resize the RDS DB instance to an instance type that has more CPU capacity.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

GOTJ 3 months, 4 weeks ago

Selected Answer: A

Option "A", but... How does an auto scaling group with min = max = 1 instances contribute to the solution?

upvoted 2 times

dhewa 9 months, 3 weeks ago

Selected Answer: A

This approach addresses both the high CPU utilization on the EC2 instance and the degraded read performance on the RDS

instance effectively.
upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: A

A sounds right
upvoted 2 times

Abbas_Abi_AWS 10 months ago

Selected Answer: A

A is correct
upvoted 2 times

ltetti 10 months ago

Selected Answer: A

Option B incorrectly suggests redirecting all read/write traffic to the replica. RDS read replicas are designed to handle read operations only, not write operations. Writes must still be handled by the primary DB instance
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #922 *Topic 1*

A company needs to grant a team of developers access to the company's AWS resources. The company must maintain a high level of security for the resources.

The company requires an access control solution that will prevent unauthorized access to the sensitive data.

Which solution will meet these requirements?

A. Share the IAM user credentials for each development team member with the rest of the team to simplify access management and to streamline development workflows.

B. Define IAM roles that have fine-grained permissions based on the principle of least privilege. Assign an IAM role to each developer. **Most Voted**

C. Create IAM access keys to grant programmatic access to AWS resources. Allow only developers to interact with AWS resources through API calls by using the access keys.

D. Create an AWS Cognito user pool. Grant developers access to AWS resources by using the user pool.

Correct Answer: *B*

Community vote distribution

B (82%) Other (18%)

Comments

LeonSauveterre 5 months ago

Selected Answer: B

A - Sharing IAM user credentials is a security risk and violates AWS best practices.
B - IAM roles allow for temporary credentials that can be automatically rotated, and that improves security.
C - Would work but managing static keys for multiple developers introduces significant operational overhead.
D - AWS Cognito is primarily designed for managing end-user authentication (for web or mobile apps or such), not for managing access to AWS resources for developers.

upvoted 1 times

hpirnaj 5 months ago

it says " team of developers " . they are not in the company's AWS account. how do you want to apply IAM roles without defining users ? I think C is the answer

upvoted 1 times

hpirnaj 5 months ago

also, roles can not assign to user/group . you need to make a policy from IAM roles first then assign policies to user/group
upvoted 1 times

KennethYY 5 months, 2 weeks ago

Selected Answer: D

A is wrong, common sense not security
B is wrong, role cannot assign to user/group
D is wrong, is designed for authentication and access control for web or mobile app users, not for internal developers accessing AWS resources.
so remain C.
upvoted 1 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: B

AWS Cognito is designed for authentication and access control for web or mobile app users, not for internal developers accessing AWS resources.
upvoted 1 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: C

AWS Cognito is designed for authentication and access control for web or mobile app users, not for internal developers accessing AWS resources.
upvoted 1 times

EllenLiu 5 months, 2 weeks ago

Sorry, go with B.
upvoted 1 times

Cpso 6 months, 1 week ago

Selected Answer: B

good practice should B. but map role to Identity center federated to corporate IDP.
upvoted 2 times

Bwhizzy 7 months, 4 weeks ago

Selected Answer: B

B is the right answer. IAM Role
upvoted 2 times

[Removed] 9 months, 3 weeks ago

B sounds right
upvoted 3 times

aragon_saa 9 months, 4 weeks ago

Selected Answer: B

Answer is B
upvoted 3 times

muhammadahmer36 9 months, 4 weeks ago

Create an AWS Cognito user pool. Grant developers access to AWS resources by using the user pool.
upvoted 1 times

mk168898 6 months, 4 weeks ago

cognito for app user auth, qns asking for access to AWS resource. your answer is wrong
upvoted 2 times

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #923 *Topic 1*

A company hosts a monolithic web application on an Amazon EC2 instance. Application users have recently reported poor performance at specific times. Analysis of Amazon CloudWatch metrics shows that CPU utilization is 100% during the periods of poor performance.

The company wants to resolve this performance issue and improve application availability.

Which combination of steps will meet these requirements MOST cost-effectively? (Choose two.)

A. Use AWS Compute Optimizer to obtain a recommendation for an instance type to scale vertically. **Most Voted**

B. Create an Amazon Machine Image (AMI) from the web server. Reference the AMI in a new launch template.

C. Create an Auto Scaling group and an Application Load Balancer to scale vertically.

D. Use AWS Compute Optimizer to obtain a recommendation for an instance type to scale horizontally.

E. Create an Auto Scaling group and an Application Load Balancer to scale horizontally. **Most Voted**

Correct Answer: *AE*

Community vote distribution

AE (50%) BE (47%) DE

Comments

rpmaws Highly Voted 8 months, 3 weeks ago

Selected Answer: AE

AWS Compute Optimizer analyzes your current EC2 instance usage and recommends the most cost-effective instance type. In this case, the current instance may not have enough CPU capacity, so scaling vertically (upgrading to a larger instance type) could provide immediate relief from the 100% CPU utilization.

upvoted 6 times

bora4motion Most Recent 3 weeks, 3 days ago

Selected Answer: BE

You can't just scale vertically every single time you have a problem. I don't understand people who voted for A and E. You can't combine A and E. You can't stop the instance to improve memory/cpu of an instance and then add horizontal scaling to it. (spend even more). B and E

upvoted 2 times

7dcef09 1 month, 3 weeks ago

Selected Answer: AE

A: Using compute optimizer helps improve performance on the monolithic web application scale up
E: Enhances availability and scalibiliy

upvoted 1 times

FlyingHawk 4 months, 1 week ago

Selected Answer: BE

The question asks for a combination of steps, meaning the selected answers should work together to form a complete solution rather than being two independent actions.
B - prepared the instance for the scaling
E - Implement Auto Scaling for performance improvement and high availability

upvoted 3 times

dfgdsfgfdgreg 4 months, 2 weeks ago

Selected Answer: BE

Since options C and D are irrelevant, let's disregard them. That leaves us with ABE. B is a prerequisite for E, so you cannot have E with B. AB will be two disjointed choices. So you have to choose B/E.

upvoted 1 times

Salilgen 5 months ago

Selected Answer: BE

The question talk about a "combination of steps" then IMO answer is BE.
BE are steps of a solution that resolve performance issue and improve application availability.
D is for improve the solution : https://docs.aws.amazon.com/compute-optimizer/latest/ug/view-asg-recommendations.html
A is a solution (not a step) but it is not the most cost effective
C is wrong: autoscaling and ALB are horizontal scaling

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: AE

What you should know:
1. Horizontal Scaling (Scaling Out/In): Adding or removing instances to distribute the load across multiple servers, improving availability and fault tolerance. Example: Using an Auto Scaling group with a load balancer to add more EC2 instances dynamically as demand increases.
2. Vertical Scaling (Scaling Up/Down): Increasing or decreasing the size (compute power, memory, etc) of a single instance, which is easier to implement for monolithic applications but has limitations (limited instance types). Example: Upgrading an instance from "t3.medium" to "m5.large".

If you can understand above, then A & E are perfect choices.
B - The performance issue is still there!
C - Auto Scaling and Load Balancers are used for horizontal scaling, not vertical scaling.
D - Horizontal scaling involves adding more instances, not changing the instance type.

upvoted 2 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: AE

it's truly difficult to answer as per the description.
there are two requirement:
1. resolve performance issue => A: use compute optimizer to scale vertically, it can provides an immediate fix for high CPU utilization
2. improve availability => B + D, B is the prerequisite for D , I prefer autoscaling if must choose only one

upvoted 1 times

Anyio 5 months, 2 weeks ago

Selected Answer: AE

The correct answers are A,E.

Option A: Correct. Using AWS Compute Optimizer to obtain a recommendation for an instance type to scale vertically is an effective first step. This service analyzes your specific application workload patterns and recommends an appropriate EC2 instance type that provides better performance for the given CPU load. This approach allows you to choose a more powerful instance that can handle peak loads more effectively without immediately shifting to a horizontal scaling architecture.

Option B: Incorrect by itself. Creating an AMI and a launch template is useful in setting up an Auto Scaling group or new instances but does not address the issue of CPU utilization or availability on its own. This process would be part of preparing for scaling but doesn't directly answer the requirement of resolving performance issues by itself.

upvoted 2 times

jfedotov 5 months, 3 weeks ago

Selected Answer: BE

horizontal scaling is a cost-effective solution

upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: DE

What about options D & E?

1 Scaling horizontally is the most cost-effective solution for handling fluctuating traffic, especially when the bottleneck is CPU utilization reaching 100% during peak times. This allows adding more instances to handle the increased load instead of upgrading to a larger instance (vertical scaling) which might be unnecessary during non-peak periods.

2 Auto Scaling group and Application Load Balancer: These are essential components for horizontal scaling, where the Auto Scaling group automatically adjusts the number of instances based on demand, and the Application Load Balancer distributes traffic evenly across available instances.

#3 AWS Compute Optimizer: This service can analyze your application usage and recommend the most cost-effective instance type for your specific needs, including suggesting options for horizontal scaling.

upvoted 1 times

Abdullah2004 7 months, 1 week ago

Selected Answer: AE

Due to monolithic application we can't scale horizontal, the scale will be vertical

upvoted 2 times

jfedotov 5 months, 3 weeks ago

you can, you'll get for ex. 2 instances with the monolithic app, I see no issues here.

upvoted 1 times

Abdullah2004 7 months, 1 week ago

I mean the instance

upvoted 1 times

XXXXXlNN 7 months, 4 weeks ago

BE makes more sense to me. Vertical scaling increases cost continuously even when the instance is in low demand. for a long run, the cost would be higher than just scale horizontally. The questions says the 'poor performance at specific times', so we just need to scale up at specific times.

Additionally, in order to scale up, we do need AMI and a launch template. Thus, a combination of B & E should be the correct answer

upvoted 2 times

Bwhizzy 7 months, 4 weeks ago

Selected Answer: AE

The answer is A and E. A for adding more CPU (Vertical scaling) and E for adding more Servers (Horizontal scaling)

upvoted 3 times

[Removed] 9 months, 2 weeks ago

AE would be the right answer

upvoted 3 times

[Removed] 9 months, 3 weeks ago

Selected Answer: BE

BE looks good

upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: BE

Answer is BE

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #924 *Topic 1*

A company runs all its business applications in the AWS Cloud. The company uses AWS Organizations to manage multiple AWS accounts.

A solutions architect needs to review all permissions that are granted to IAM users to determine which IAM users have more permissions than required.

Which solution will meet these requirements with the LEAST administrative overhead?

A. Use Network Access Analyzer to review all access permissions in the company's AWS accounts.

B. Create an AWS CloudWatch alarm that activates when an IAM user creates or modifies resources in an AWS account.

C. Use AWS Identity and Access Management (IAM) Access Analyzer to review all the company's resources and accounts. **Most Voted**

D. Use Amazon Inspector to find vulnerabilities in existing IAM policies.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

muhammadahmer36 9 months, 4 weeks ago

Selected Answer: C

https://aws.amazon.com/iam/access-analyzer/

upvoted 2 times

swati1508 10 months ago

C is correct

upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: C

Answer is C

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #925 *Topic 1*

A company needs to implement a new data retention policy for regulatory compliance. As part of this policy, sensitive documents that are stored in an Amazon S3 bucket must be protected from deletion or modification for a fixed period of time.

Which solution will meet these requirements?

A. Activate S3 Object Lock on the required objects and enable governance mode.

B. Activate S3 Object Lock on the required objects and enable compliance mode. Most Voted

C. Enable versioning on the S3 bucket. Set a lifecycle policy to delete the objects after a specified period.

D. Configure an S3 Lifecycle policy to transition objects to S3 Glacier Flexible Retrieval for the retention duration.

Correct Answer: *B*

Community vote distribution

B (89%) C (11%)

Comments

Sarayounisaldossary 2 months ago

Selected Answer: B

This is tricky between A and C but-- the question focused on compliance and FIXED Period of time... governance feature will let certain users to override B

upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: B

S3 Object Lock:
- This feature allows you to lock objects in an S3 bucket, preventing them from being deleted or modified for a specified retention period.

Compliance mode:
- When enabled, S3 Object Lock in compliance mode provides the most stringent protection, ensuring that once an object is locked, even the account root user cannot delete or modify it, making it ideal for regulatory compliance scenarios where data must be preserved for a fixed period

upvoted 3 times

SmartGil2024 6 months, 2 weeks ago

Selected Answer: C

The question says : : "for a fixed period of time.", why not C ? Because the items can be deleted after this fix period of time./ with compliance mode - no one can delete, it is true, but we need to also pay attention to the fixed period of time.

upvoted 1 times

JA2018 6 months, 1 week ago

While versioning helps maintain historical data, a lifecycle policy alone cannot guarantee that sensitive documents will not be deleted prematurely, as it can be modified or deleted.

upvoted 1 times

[Removed] 9 months, 3 weeks ago

Selected Answer: B

B is correct

upvoted 2 times

komorebi 10 months ago

Selected Answer: B

Answer is B

upvoted 2 times

swati1508 10 months ago

B compliance mode - no one can delete

upvoted 3 times

ccceb01 9 months, 1 week ago

Except god. lol

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #926 *Topic 1*

A company runs its customer-facing web application on containers. The workload uses Amazon Elastic Container Service (Amazon ECS) on AWS Fargate. The web application is resource intensive.

The web application needs to be available 24 hours a day, 7 days a week for customers. The company expects the application to experience short bursts of high traffic. The workload must be highly available.

Which solution will meet these requirements MOST cost-effectively?

A. Configure an ECS capacity provider with Fargate. Conduct load testing by using a third-party tool. Rightsize the Fargate tasks in Amazon CloudWatch.

B. Configure an ECS capacity provider with Fargate for steady state and Fargate Spot for burst traffic.

C. Configure an ECS capacity provider with Fargate Spot for steady state and Fargate for burst traffic.

D. Configure an ECS capacity provider with Fargate. Use AWS Compute Optimizer to rightsize the Fargate task. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (55%) B (45%)

Comments

rosanna Highly Voted 6 months, 1 week ago

Selected Answer: D

spots might be interrupted and the questions says it has to be available 24/7, rightsizing addresses the cost optimization point but the spot is so risk so I'll go for D

upvoted 5 times

FlyingHawk 4 months, 3 weeks ago

If it is a stateless and fault tolerate web app, it should be okay to be interrupted.
https://repost.aws/knowledge-center/fargate-spot-termination-notice

upvoted 2 times

7dcef09 Most Recent 1 month, 3 weeks ago

Selected Answer: B

B is a better answer
upvoted 2 times

Sarayounisaldossary 2 months ago

Selected Answer: B

Use Fargate Spot for burst traffic — it’s significantly cheaper (up to 70% savings)

ECS Capacity Providers let you define a mix strategy for both types of compute

This solution provides: High availability via regular Fargate- Cost savings (via Spot)- Scalability for traffic bursts without manual management

why other options are incorrect?
A.
Only uses standard Fargate with performance tuning tools.

Does not address cost optimization or scaling for traffic bursts.

No use of Fargate Spot, which misses out on potential savings.

C.
Using Fargate Spot for steady-state workloads is not reliable.

Spot tasks can be interrupted at any time, which is not acceptable for a 24/7 high-availability app.

D.
AWS Compute Optimizer helps with right-sizing but doesn't handle scaling or burst traffic.
upvoted 1 times

dfgdsfgfdgreg 4 months, 2 weeks ago

Selected Answer: B

Compute Optimizer can not optimise fargate tasks.
upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: B

https://repost.aws/knowledge-center/fargate-spot-termination-notice
B is correct as you can use ECS task draining to gracefully handle Spot interruptions.
upvoted 1 times

BugsyWarribwoy 5 months ago

Selected Answer: B

B is the valid answer. Use Fargate for steady-state tasks to ensure high availability and reliability, and Fargate Spot for burst tasks to handle high traffic at a reduced cost. This combination meets the requirements most cost-effectively.
D is wrong because rightsizing alone does not optimize for burst traffic scenarios or reduce costs as effectively as incorporating Fargate Spot.
upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: B

A - It doesn't address the cost optimization for burst traffic. Rightsizing ensures you're not over-provisioning your baseline resources, but it doesn't provide a mechanism to dynamically scale up with cost-effective Spot instances during traffic spikes.
B - The application can scale up quickly with Fargate Spot tasks to handle the increased load, and if a Spot task is interrupted, the remaining Fargate tasks can continue to serve traffic (and the system can try to acquire new Spot capacity).
C - Fargate Spot is not recommended for steady-state workloads because it can be interrupted. This would compromise the application’s 24/7 availability.
D - Like A, it fails to leverage Fargate Spot for cost savings.

Combining Fargate for steady state with Fargate Spot for burst traffic, the 24/7 HA is actually guaranteed.
upvoted 3 times

Liongeek 5 months, 2 weeks ago

Selected Answer: D

It's D.

Compute Optimizer supports recommendations for Fargate Tasks
https://aws.amazon.com/about-aws/whats-new/2022/12/aws-compute-optimizer-amazon-ecs-services-aws-fargate/

Option B and C recommend Spot, but that's not an option since the the question asks for 24x7 availability.
upvoted 3 times

thiahthura 5 months, 2 weeks ago

Selected Answer: D

D is right answer. Company desire is 24 * 7. So I think spot should be.
upvoted 2 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: D

available for 24 * 7, spot is out as it can be interrupted any time.
upvoted 4 times

spoved 8 months, 2 weeks ago

Selected Answer: D

https://aws.amazon.com/blogs/aws-cloud-financial-management/how-to-take-advantage-of-rightsizing-recommendation-preferences-in-compute-optimizer/
be available 24 hours a day, 7 days a week
must be highly available
=> D
upvoted 4 times

FlyingHawk 4 months, 3 weeks ago

Rightsizing is important for cost optimization, but it does not address the need for handling burst traffic.
upvoted 1 times

dhewa 9 months, 3 weeks ago

Selected Answer: B

This combination leverages the cost benefits of Fargate Spot for burst traffic while ensuring steady performance with regular Fargate instances.
upvoted 2 times

pujithacg8 10 months ago

B is the right answer
upvoted 3 times

swati1508 10 months ago

B- for short work use spot
upvoted 3 times

flaviobrf 10 months, 1 week ago

Selected Answer: B

B is the right choice, the application must be available 24/7
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #927 *Topic 1*

A company is building an application in the AWS Cloud. The application is hosted on Amazon EC2 instances behind an Application Load Balancer (ALB). The company uses Amazon Route 53 for the DNS.

The company needs a managed solution with proactive engagement to detect against DDoS attacks.

Which solution will meet these requirements?

A. Enable AWS Config. Configure an AWS Config managed rule that detects DDoS attacks.

B. Enable AWS WAF on the ALCreate an AWS WAF web ACL with rules to detect and prevent DDoS attacks. Associate the web ACL with the ALB.

C. Store the ALB access logs in an Amazon S3 bucket. Configure Amazon GuardDuty to detect and take automated preventative actions for DDoS attacks.

D. Subscribe to AWS Shield Advanced. Configure hosted zones in Route 53. Add ALB resources as protected resources.
Most Voted

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

dhewa Highly Voted 9 months, 3 weeks ago

Selected Answer: D

Tip: DDoS = Shield, SQL injection = WAF

upvoted 5 times

LeonSauveterre Most Recent 5 months ago

Selected Answer: D

For your references:
1. AWS WAF (Web Application Firewall) - SQL injections, cross-site scripting (XSS), and other common web exploits.
2. Shield Standard - common DDoS attacks; Shield Advanced - Provides advanced DDoS protection with 24/7 support, cost protection for scaling due to attacks, and more detailed attack analytics.
3. AWS Network Firewall - Blocks unauthorized access and inspect traffic for VPCs.

4. Amazon GuardDuty - Detects threats by analyzing AWS logs like CloudTrail, DNS, and VPC Flow Logs. Identifies potential malicious activities like brute force attacks and data exfiltration.
5. AWS Macie - Helps identify and protect sensitive data like PII stored in S3 buckets.

upvoted 2 times

LeonSauveterre 5 months ago

6. AWS Detective - Analyzes and visualizes security-related data to identify root causes of potential security issues.
7. AWS Config - Monitors and evaluates configurations of AWS resources to ensure compliance with security best practices.
8. AWS Inspector - Automatically scans EC2 instances and container-based workloads for known vulnerabilities.
9. AWS CloudTrail - Tracks API calls and user activity for auditing and compliance.

upvoted 2 times

AMEJack 6 months, 1 week ago

Selected Answer: D

Shield Advanced signals a DDoS detection and starts analyzing the traffic for an attack signature. If a signature is found, it is first tested on past traffic to reduce the risk of false positive, then if it's safe to use it, a corresponding WAF rule is placed in the previously created rule group. After a certain time, when the attack is stopped, the rule is automatically removed from the rule group. When successful, this process takes several minutes.

upvoted 2 times

komorebi 10 months ago

Selected Answer: D

Answer is D

upvoted 2 times

swati1508 10 months ago

D for DDOS shield advance

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #928 *Topic 1*

A company hosts a video streaming web application in a VPC. The company uses a Network Load Balancer (NLB) to handle TCP traffic for real-time data processing. There have been unauthorized attempts to access the application.

The company wants to improve application security with minimal architectural change to prevent unauthorized attempts to access the application.

Which solution will meet these requirements?

A. Implement a series of AWS WAF rules directly on the NLB to filter out unauthorized traffic.

B. Recreate the NLB with a security group to allow only trusted IP addresses. **Most Voted**

C. Deploy a second NLB in parallel with the existing NLB configured with a strict IP address allow list.

D. Use AWS Shield Advanced to provide enhanced DDoS protection and prevent unauthorized access attempts.

Correct Answer: *B*

Community vote distribution

B (52%) D (45%) A

Comments

bora4motion 3 weeks, 2 days ago

Selected Answer: D

It can't be B. You are "netflix"lets say. You don't know which IP addresses are trusted to use them in SG. It's stupid.

upvoted 1 times

GOTJ 3 months, 4 weeks ago

Selected Answer: B

Several links below prove that option "B" is correct, regardless of the "NLB recreation" drawback.

It's time to kick option C out (it was my first guess, I must confess)

DDoS prevention, even though useful, is not required by the company. The company is interested SOLELY in prevent "unauthorized access" to the application. In the absence of more specific information, I came to the conclusion that the company detected some "weird ips" trying to access the application. After some investigation, I came to the conclusion that AWS Shield Advanced doesn't provide this feature for ips if they weren't previously involved in a DDoS attack.

In an on-premises environment, I would setup a firewall rule with an allow list to accomplish this requirement, which is precisely what option "B" offers.
upvoted 2 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: B

https://aws.amazon.com/blogs/containers/network-load-balancers-now-support-security-groups/
upvoted 4 times

LeonSauveterre 5 months ago

Selected Answer: D

A - AWS WAF works at Layer 7 (application layer) and is designed for HTTP/HTTPS traffic, so WAF works only with Application Load Balancers (ALB), API Gateway, and CloudFront.
C - Works but comes with unnecessary complexity and introduces architectural changes.
So A & C are out. I'm actually torn between B and D.

If the primary issue is unauthorized IP access and not large-scale DDoS attacks, then B might be the answer, but NLBs do not directly associate with security groups. Security groups are applied to the targets of the NLB (EC2 instances, IP addrs, or ALBs). Also, this is an architectural change.

On the other hand, if the primary concern includes DDoS attacks, option D (AWS Shield Advanced) becomes more relevant but this is so much more expensive and may still be overkill for simple IP-based access control. I'm gonna go with D if it shows up in my exam.
upvoted 3 times

FlyingHawk 4 months, 3 weeks ago

AWS NLB supports the security group now: https://aws.amazon.com/blogs/containers/network-load-balancers-now-support-security-groups/
upvoted 3 times

Salilgen 5 months ago

https://docs.aws.amazon.com/elasticloadbalancing/latest/network/load-balancer-security-groups.html
upvoted 2 times

Denise123 5 months, 1 week ago

Selected Answer: B

Tricky one. About option A > AWS WAF does not support Network Load Balancers (NLBs) directly. NLBs operate at the transport layer (Layer 4), while AWS WAF is designed to work with Application Load Balancers (ALBs) at the application layer (Layer 7).....Given that the requirement is to improve application security for a Network Load Balancer with minimal architectural changes, the most appropriate solution would be Option B
upvoted 2 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: B

shield is only for DDos attacks protection
upvoted 1 times

Anyio 5 months, 2 weeks ago

Selected Answer: D

The correct answer is D. Use AWS Shield Advanced to provide enhanced DDoS protection and prevent unauthorized access attempts.

Explanation:

Option D: Correct. AWS Shield Advanced is designed to protect against DDoS attacks, which can be a source of unauthorized access attempts. It provides enhanced protection features for applications behind a Network Load Balancer, offering additional security measures without requiring significant architectural changes. By leveraging AWS Shield Advanced, the company can gain comprehensive DDoS protection tailored for use with their existing NLB setup.

Option B: Incorrect. NLBs do not support security groups which are applicable to instances, not to the NLB itself. In addition, recreating the NLB to deal with unauthorized access attempts does not align with the requirement for minimal architectural

change.
upvoted 3 times

dragossky 5 months, 3 weeks ago

Selected Answer: D

Use AWS Shield Advanced to provide enhanced DDoS protection and prevent unauthorized access attempts.
upvoted 4 times

ckhemani 6 months, 1 week ago

Selected Answer: B

Real-time data processing normally use RTP Protocol which uses a range of ports to deliver audio and video streams. It doesn't specifically says HTTPS so i assume, it can't use WAF to control the traffic since it operates in HTTP/HTTPS Level only.

Not designed for real-time:
HTTP is primarily designed for request-response communication, which involves sending a request and then waiting for a full response, making it less efficient for continuous data streams needed in real-time applications
upvoted 1 times

tttttttttttttttttttttt 6 months, 2 weeks ago

Selected Answer: A

Why is it not A?
upvoted 1 times

ARV14 6 months, 1 week ago

Waf supports ALB layer7, not nlb
upvoted 2 times

Sergantus 6 months, 3 weeks ago

Selected Answer: D

The answer should be D. It makes no sense to pick B for a public app in cases of DDoS, SGs wouldn't help with that. It's like, the closer the questions end, the more trolls left.
upvoted 3 times

mk168898 6 months, 4 weeks ago

I don't think B is correct. if you only allow selected IPs to access then this company cannot host their video streaming service to the public.

D should be the correct answer. AWS shield advanced if I rmb correctly prevent unauthorised attempts
upvoted 1 times

Jeyaluxshan 9 months ago

Network Load Balancers (NLB) now supports security groups, enabling you to filter the traffic that your NLB accepts and forwards to your application. Using security groups, you can configure rules to help ensure that your NLB only accepts traffic from trusted IP addresses, and centrally enforce access control policies. This improves your application's security posture and simplifies operations
upvoted 3 times

AbhiBK 9 months, 1 week ago

Answer is D
upvoted 3 times

[Removed] 9 months, 3 weeks ago

Selected Answer: B

B is correct
upvoted 1 times

komorebi 10 months, 1 week ago

Selected Answer: B

Answer is B

Answer is B
upvoted 1 times

example_ 10 months, 1 week ago

Selected Answer: B

https://aws.amazon.com/about-aws/whats-new/2023/08/network-load-balancer-supports-security-groups/
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #929 *Topic 1*

A healthcare company is developing an AWS Lambda function that publishes notifications to an encrypted Amazon Simple Notification Service (Amazon SNS) topic. The notifications contain protected health information (PHI).

The SNS topic uses AWS Key Management Service (AWS KMS) customer managed keys for encryption. The company must ensure that the application has the necessary permissions to publish messages securely to the SNS topic.

Which combination of steps will meet these requirements? (Choose three.)

A. Create a resource policy for the SNS topic that allows the Lambda function to publish messages to the topic. **Most Voted**

B. Use server-side encryption with AWS KMS keys (SSE-KMS) for the SNS topic instead of customer managed keys.

C. Create a resource policy for the encryption key that the SNS topic uses that has the necessary AWS KMS permissions. **Most Voted**

D. Specify the Lambda function's Amazon Resource Name (ARN) in the SNS topic's resource policy.

E. Associate an Amazon API Gateway HTTP API with the SNS topic to control access to the topic by using API Gateway resource policies.

F. Configure a Lambda execution role that has the necessary IAM permissions to use a customer managed key in AWS KMS. **Most Voted**

Correct Answer: *ACF*

Community vote distribution

ACF (70%) ADF (30%)

Comments

Abbas_Abi_AWS Highly Voted 10 months ago

Selected Answer: ACF

A C F is corrcet

upvoted 8 times

GOTJ Most Recent 3 months, 4 weeks ago

Selected Answer: ADF

Does a SNS topic need to decrypt an encrypted message in order to deliver it to its subscribers? If so, option "C" should be in. If not, it should be out. And I think is out, unless SNS topic needs to apply certain logic to the message prior to subscription deliver (filters).

This question requires a CSE-KMS approach, which force the Lambda function to encrypt the message before publishing, so option "F" is a must (I think everybody here is on the same page with this one).

Also, option "A" should be obvious by now. And even though option "D" is redundant, it's technically correct, so it's a better choice.

upvoted 2 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: ACF

For C - https://docs.aws.amazon.com/sns/latest/dg/sns-enable-encryption-for-topic.html
Permissions for custom KMS keys – If using a custom KMS key, include the following in the key policy to allow Amazon SNS to encrypt and decrypt messages:

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

For A: https://repost.aws/knowledge-center/sns-topic-lambda and https://repost.aws/knowledge-center/sns-topics-iam-errors-subscriber
For F: https://repost.aws/knowledge-center/lambda-kmsaccessdeniedexception-errors

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: ACF

A - This is a must-do.
B - As question stated, we have no reason to replace the keys.
C - This ensures the function can encrypt messages when publishing.
D - While specifying the Lambda function's ARN in the SNS topic's resource policy is part of the solution (it's how you grant the "sns:Publish" permission), it's not sufficient on its own. You also need the KMS key policy and the Lambda execution role permissions. This option is a part of option A, not a standalone answer.
E - The Lambda function interacts directly with the SNS topic, and API Gateway is not relevant.
F - Yes, so that this role will have policies that allow the "kms:Encrypt", "kms:GenerateDataKey," and "kms:Decrypt" actions on the specific KMS key used to encrypt the topic.

upvoted 3 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: ACF

I don't understand why lambda needs KMS to encrypt message.
In this scenario, the Lambda function does not need direct permissions to use the KMS key because the encryption and decryption are fully managed by Amazon SNS as part of the SSE-KMS feature.
SNS is the only service lambda will talk to, The Lambda interacts with SNS using HTTPS and does not directly deal with the encrypted data or the KMS key. so Lambda only needs the permission to publish messages to the SNS topic.
I would like to choose F only if #F answer can be updated as below:
F. Configure a Lambda execution role that has the necessary IAM permissions to publish to the SNS topic.
A: Create a resource policy for the SNS topic => grants lambda the ability to publish messages to the topic
C: The KMS key resource policy must allow SNS to use the key for encryption and decryption.
F: Configure the Lambda Execution Role => SNS permissions to publish to the topic

upvoted 2 times

JA2018 6 months, 1 week ago

Selected Answer: ADF

#A: This is essential because the resource policy on the SNS topic will define which entities (like the Lambda function) are allowed to publish messages to it.

#D: By specifying the Lambda function's ARN in the SNS topic policy, you clearly grant access only to that specific Lambda function.

#F: Since the SNS topic uses a customer-managed KMS key, the Lambda execution role must have the necessary permissions to use that key for encryption/decryption when publishing messages.

upvoted 2 times

FlyingHawk 4 months, 3 weeks ago

D is redundant because Option A already covers the need to allow the Lambda function to publish to the SNS topic. The ARN of the Lambda function would be included in the resource policy created in Option A.

upvoted 1 times

JA2018 6 months, 1 week ago

Why the other options are not correct:

#B: The question already states that the SNS topic is using customer-managed KMS keys, so there's no need to switch to server-side encryption with AWS managed keys.

#C: While technically you could create a resource policy on the encryption key itself, it's not the most secure approach as it would grant access to the key itself, not just the ability to use it for SNS encryption.

#E: Introducing an API Gateway layer is unnecessary complexity for this scenario, as you can directly control access to the SNS topic using the Lambda execution role and its permissions.

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

C is correct since the SNS topic uses a customer managed key for encryption. The key's resource policy must allow the SNS service to use the key for encryption and decryption.

upvoted 1 times

JA2018 6 months, 1 week ago

Key takeaway:

- To ensure the Lambda function can securely publish messages to an encrypted SNS topic, you need to properly configure the SNS topic resource policy to allow the Lambda function access and make sure the Lambda execution role has the necessary KMS permissions to use the customer-managed encryption key.

upvoted 1 times

elmyth 7 months, 1 week ago

Selected Answer: ACF

D is correct too and C is not clear, but seems like it is about KMS policy and adding permissions for sns service which has to be added in case of CMK

upvoted 4 times

agbor_tambe 8 months, 1 week ago

Selected Answer: ADF

my answer

upvoted 1 times

Sergantus 6 months, 3 weeks ago

D is like a part of A, so it makes no sense to pick both, it should be A C F.

upvoted 1 times

progounick 9 months, 1 week ago

Selected Answer: ACF

ChatGPT agrees with me

upvoted 3 times

komorebi 10 months, 1 week ago

Selected Answer: ADF

Answer is ADF

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #930 *Topic 1*

A company has an employee web portal. Employees log in to the portal to view payroll details. The company is developing a new system to give employees the ability to upload scanned documents for reimbursement. The company runs a program to extract text-based data from the documents and attach the extracted information to each employee's reimbursement IDs for processing.

The employee web portal requires 100% uptime. The document extract program runs infrequently throughout the day on an on-demand basis. The company wants to build a scalable and cost-effective new system that will require minimal changes to the existing web portal. The company does not want to make any code changes.

Which solution will meet these requirements with the LEAST implementation effort?

A. Run Amazon EC2 On-Demand Instances in an Auto Scaling group for the web portal. Use an AWS Lambda function to run the document extract program. Invoke the Lambda function when an employee uploads a new reimbursement document. Most Voted

B. Run Amazon EC2 Spot Instances in an Auto Scaling group for the web portal. Run the document extract program on EC2 Spot Instances. Start document extract program instances when an employee uploads a new reimbursement document.

C. Purchase a Savings Plan to run the web portal and the document extract program. Run the web portal and the document extract program in an Auto Scaling group.

D. Create an Amazon S3 bucket to host the web portal. Use Amazon API Gateway and an AWS Lambda function for the existing functionalities. Use the Lambda function to run the document extract program. Invoke the Lambda function when the API that is associated with a new document upload is called.

Correct Answer: *A*

Community vote distribution

A (58%) C (33%) D (8%)

Comments

hpirnaj Highly Voted 5 months ago

Selected Answer: C

i hate these types of questions . there is no perfect answer
A:on demand "not cost-effective " . also , how do you want to invoke Lambda , there is no API call

A:on demand "not cost effective" , also , how do you want to invoke Lambda . there is no API call
B: spot instance ("web portal requires 100% uptime")
D:s3 is for static website . it can work but payroll needs to update every 2 weeks and require extensive code change
C in my opinion is the only correct answer. it is scalable "ASG", no code change and cost effective "saving plan"
upvoted 6 times

7dcef09 Most Recent 1 month, 3 weeks ago

Selected Answer: C

is this supposed to be a confusing question? C is scalable and most cost effective.
upvoted 1 times

Salilgen 5 months ago

Selected Answer: C

Spot Instances are not suitable for workloads requiring 100% uptime then B is out.
D is the best solution but it requires some implementation effort.
Savings plan (that include Lambda too...) is cheaper than on-demand (option A).
Moreover, C doesn't require any code change or other AWS services.
upvoted 1 times

tch 2 months, 4 weeks ago

which saving plan to purchase ? Compute Savings Plans, EC2 Instance Savings Plans, and Amazon SageMaker Savings Plans???
upvoted 1 times

ckhemani 6 months, 1 week ago

Selected Answer: A

Answer: A
This solution offers the most scalable and cost-effective approach with minimal changes to the existing web portal and no code modifications.

Why Not Other Options?:

Option B (Spot Instances): Spot Instances are not suitable for workloads requiring 100% uptime, as they can be terminated by AWS with short notice.

Option C (Savings Plan): A Savings Plan could reduce costs but does not address the requirement for running the document extraction program efficiently or without code changes.

Option D (S3 with API Gateway and Lambda): This would require significant changes to the existing web portal setup, including moving the portal to S3 and reconfiguring its architecture, which contradicts the requirement of minimal implementation effort and no code change.
upvoted 3 times

XXXXXlNN 7 months, 4 weeks ago

D makes more sense if the company asks to redesign the whole thing to achieve better operational management, performance, cost effective, etc. However, it requires us to provide solution with MINIMUM change... thus A it is I guess.
upvoted 3 times

JoeTromundo 8 months, 2 weeks ago

Selected Answer: A

"The company does not want to make any code changes."
Option D requires a complete re-architecture of the web portal to be hosted on Amazon S3 and API Gateway, which involves significant changes to the existing system. This does not align with the requirement of minimal changes to the current setup.
upvoted 4 times

rpmaws 8 months, 3 weeks ago

Selected Answer: A

because they don't want any change in code so A is correct.
upvoted 3 times

progounick 9 months, 1 week ago

Selected Answer: A

A ChatGPT agrees with me
upvoted 2 times

progounick 9 months, 2 weeks ago

Selected Answer: D

I think D is better choice. Even though A makes sense too, D seems the correct one
upvoted 1 times

RealPro111 9 months, 2 weeks ago

Selected Answer: D

Least effort, webportal is simply an interface: D
upvoted 1 times

[Removed] 9 months, 3 weeks ago

Selected Answer: A

A sounds right
upvoted 2 times

744fdad 9 months, 4 weeks ago

i think D
upvoted 2 times

komorebi 10 months ago

Answer is A
upvoted 1 times

swati1508 10 months ago

I thinks it's D
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #931 *Topic 1*

A media company has a multi-account AWS environment in the us-east-1 Region. The company has an Amazon Simple Notification Service (Amazon SNS) topic in a production account that publishes performance metrics. The company has an AWS Lambda function in an administrator account to process and analyze log data.

The Lambda function that is in the administrator account must be invoked by messages from the SNS topic that is in the production account when significant metrics are reported.

Which combination of steps will meet these requirements? (Choose two.)

A. Create an IAM resource policy for the Lambda function that allows Amazon SNS to invoke the function. **Most Voted**

B. Implement an Amazon Simple Queue Service (Amazon SQS) queue in the administrator account to buffer messages from the SNS topic that is in the production account. Configure the SQS queue to invoke the Lambda function.

C. Create an IAM policy for the SNS topic that allows the Lambda function to subscribe to the topic. **Most Voted**

D. Use an Amazon EventBridge rule in the production account to capture the SNS topic notifications. Configure the EventBridge rule to forward notifications to the Lambda function that is in the administrator account.

E. Store performance metrics in an Amazon S3 bucket in the production account. Use Amazon Athena to analyze the metrics from the administrator account.

Correct Answer: *AC*

Community vote distribution

AC (52%) | AB (33%) | Other (15%)

Comments

siheom Highly Voted 10 months ago

Selected Answer: AB

VOTE A,B

upvoted 5 times

Anyio Most Recent 5 months, 2 weeks ago

Selected Answer: AC

The Correct answer is A C

The Correct answer is A,C.

Using the Amazon SNS console, add a cross-account AWS Lambda subscription to an Amazon SNS topic.

the Lambda function resource policy allows SNS to invoke the function.
the SNS topic access policy allows Lambda to subscribe to the topic.
Note: The SNS topic resides in account A and the Lambda function resides in account B.

upvoted 1 times

Anyio 5 months, 2 weeks ago

Selected Answer: AB

The correct answers are A,B.

Explanation:

Option A: Correct. Creating an IAM resource policy for the Lambda function that allows Amazon SNS to invoke the function is necessary for SNS to have permission to trigger Lambda. This policy ensures that the Lambda function can be invoked by a service principal from the SNS service.

Option B: Correct. Using an Amazon Simple Queue Service (Amazon SQS) queue as an intermediary buffer allows for decoupling the SNS topic from the Lambda function, providing more reliability and handling burst traffic effectively. In this setup, the SNS topic can publish to the SQS queue, and the queue can then trigger the Lambda function to process the messages.

upvoted 2 times

Anyio 5 months, 2 weeks ago

Sorry this is the wrong answer (or second best answer) :). The Correct answer is A,C.
Using the Amazon SNS console, add a cross-account AWS Lambda subscription to an Amazon SNS topic.

the Lambda function resource policy allows SNS to invoke the function.
the SNS topic access policy allows Lambda to subscribe to the topic.
Note: The SNS topic resides in account A and the Lambda function resides in account B.

upvoted 2 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: AC

A: resource-policy for Lambda: should grant SNS to access lambda permission
C: resource-policy for SNS: should specify who can subscribe SNS topic

upvoted 3 times

JA2018 6 months, 1 week ago

Selected Answer: AD

#A: This is the most direct way to allow the SNS topic in the production account to trigger the Lambda function in the administrator account. By creating an IAM policy on the Lambda function that grants SNS permission to invoke it, you establish the necessary access control.

#D: Using an EventBridge rule in the production account allows you to filter and route the SNS notifications specifically to the Lambda function in the administrator account, providing greater control and flexibility over the event delivery .

upvoted 1 times

agbor_tambe 8 months, 1 week ago

Selected Answer: AC

most reasonable

upvoted 1 times

mooondooo 8 months, 2 weeks ago

Selected Answer: AC

Probably A and C

https://repost.aws/knowledge-center/sns-with-crossaccount-lambda-subscription

upvoted 3 times

progounick 9 months, 2 weeks ago

Selected Answer: AC

Selected Answer: AC

A and C seem to be the best answer

upvoted 1 times

dhewa 9 months, 3 weeks ago

Selected Answer: AC

No need to complicate stuff, AWS services already exist only permissions are missing. A&C will set up the necessary permissions and subscriptions for cross-account invocation of the Lambda function by the SNS topic.

upvoted 2 times

523db89 9 months, 3 weeks ago

A,C correct - While using SQS could be a solution for buffering messages, it introduces additional complexity

upvoted 1 times

jamesukae 9 months, 3 weeks ago

Selected Answer: BE

For me AB is contradict , why we invoke lambda function by both SNS and SQS?

I think BE is correct answer because question also need solution to analyze data.

upvoted 2 times

nebajp 10 months ago

correct answer is AD

upvoted 4 times

JA2018 6 months, 1 week ago

- #A: This is the most direct way to allow the SNS topic in the production account to trigger the Lambda function in the administrator account. By creating an IAM policy on the Lambda function that grants SNS permission to invoke it, you establish the necessary access control [A, D].
- #D: Using an EventBridge rule in the production account allows you to filter and route the SNS notifications specifically to the Lambda function in the administrator account, providing greater control and flexibility over the event delivery [D].

upvoted 2 times

GOTJ 4 months ago

I like your reasoning for option "D". However, with this setup, EventBridge should be the service that invoke the Lambda function:

Perf Metrics --> SNS --> EventBridge (for filtering and routing "the significant metrics") --> Lambda

Since you've detached Lambda from SNS adding a service in between, SNS should no longer invoke the Lambda Function and option "A" would be wrong, isn't it?

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #932 *Topic 1*

A company is migrating an application from an on-premises location to Amazon Elastic Kubernetes Service (Amazon EKS). The company must use a custom subnet for pods that are in the company's VPC to comply with requirements. The company also needs to ensure that the pods can communicate securely within the pods' VPC.

Which solution will meet these requirements?

A. Configure AWS Transit Gateway to directly manage custom subnet configurations for the pods in Amazon EKS.

B. Create an AWS Direct Connect connection from the company's on-premises IP address ranges to the EKS pods.

C. Use the Amazon VPC CNI plugin for Kubernetes. Define custom subnets in the VPC cluster for the pods to use. **Most Voted**

D. Implement a Kubernetes network policy that has pod anti-affinity rules to restrict pod placement to specific nodes that are within custom subnets.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

officedepotadmin Highly Voted 10 months ago

Selected Answer: C

The Amazon VPC Container Network Interface (CNI) plugin is the default network plugin for Amazon EKS. It allows Kubernetes pods to receive IP addresses from a VPC's subnet and enables pods to communicate securely within the VPC as if they were native VPC resources.

upvoted 8 times

Sarayounisaldossary Most Recent 2 months ago

Selected Answer: C

The Amazon VPC CNI plugin is a Kubernetes networking plugin for EKS.
where it provides: a unique VPC IP address for each pod and native integration of Kubernetes networking into AWS VPC

- Pods can communicate like any EC2 instance on the same subnet.

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: C

A - AWS Transit Gateway is used to manage connectivity between multiple VPCs, on-premises networks, or other AWS services. It does not manage pod-level subnet configurations in Amazon EKS.
B - This can't address the specific requirement to configure custom subnets for pods within the VPC.
C - This is the recommended approach.
D - Pod anti-affinity rules are used for scheduling constraints, not for subnet allocation or secure communication within the VPC.

upvoted 1 times

mooondooo 8 months, 2 weeks ago

Selected Answer: C

Probably C

https://repost.aws/knowledge-center/eks-custom-subnet-for-pod

upvoted 3 times

dhewa 9 months, 3 weeks ago

Selected Answer: C

C all the way.

upvoted 2 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #933 *Topic 1*

A company hosts an ecommerce application that stores all data in a single Amazon RDS for MySQL DB instance that is fully managed by AWS. The company needs to mitigate the risk of a single point of failure.

Which solution will meet these requirements with the LEAST implementation effort?

A. Modify the RDS DB instance to use a Multi-AZ deployment. Apply the changes during the next maintenance window. **Most Voted**

B. Migrate the current database to a new Amazon DynamoDB Multi-AZ deployment. Use AWS Database Migration Service (AWS DMS) with a heterogeneous migration strategy to migrate the current RDS DB instance to DynamoDB tables.

C. Create a new RDS DB instance in a Multi-AZ deployment. Manually restore the data from the existing RDS DB instance from the most recent snapshot.

D. Configure the DB instance in an Amazon EC2 Auto Scaling group with a minimum group size of three. Use Amazon Route 53 simple routing to distribute requests to all DB instances.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

example_ Highly Voted 10 months, 1 week ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_UpgradeDBInstance.Maintenance.html

upvoted 5 times

2f2127f Most Recent 2 weeks, 5 days ago

Selected Answer: A

It's because I said so

upvoted 1 times

Changwha 6 months, 2 weeks ago

A: This solution provides automated failover and high availability with minimal implementation effort.

upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: A

Answer is A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #934 *Topic 1*

A company has multiple Microsoft Windows SMB file servers and Linux NFS file servers for file sharing in an on-premises environment. As part of the company's AWS migration plan, the company wants to consolidate the file servers in the AWS Cloud.

The company needs a managed AWS storage service that supports both NFS and SMB access. The solution must be able to share between protocols. The solution must have redundancy at the Availability Zone level.

Which solution will meet these requirements?

A. Use Amazon FSx for NetApp ONTAP for storage. Configure multi-protocol access. **Most Voted**

B. Create two Amazon EC2 instances. Use one EC2 instance for Windows SMB file server access and one EC2 instance for Linux NFS file server access.

C. Use Amazon FSx for NetApp ONTAP for SMB access. Use Amazon FSx for Lustre for NFS access.

D. Use Amazon S3 storage. Access Amazon S3 through an Amazon S3 File Gateway.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

FlyingHawk 4 months, 3 weeks ago

Selected Answer: A

https://aws.amazon.com/fsx/netapp-ontap/features/
Key Features of Amazon FSx for NetApp ONTAP:
Multi-Protocol Support: Simultaneously supports SMB, NFS, and iSCSI protocols.

Multi-AZ Redundancy: Provides high availability and automatic failover across Availability Zones.

Scalability: Scales storage capacity and performance independently.

Data Management: Includes advanced features like snapshots, cloning, and data deduplication.

Integration: Seamlessly integrates with AWS services and on-premises environments.

upvoted 2 times

upvoted 2 times

Jeyaluxshan 9 months ago

Amazon FSx for NetApp ONTAP uses Single-AZ and Multi-AZ deployment types. You can choose from four options: Single-AZ 1, Single-AZ 2, Multi-AZ 1, and Multi-AZ 2

upvoted 2 times

ccceb01 9 months, 1 week ago

Selected Answer: A

https://aws.amazon.com/fsx/netapp-ontap/

upvoted 3 times

komorebi 10 months, 1 week ago

Selected Answer: A

Answer is A

upvoted 2 times

example_ 10 months, 1 week ago

Selected Answer: A

https://aws.amazon.com/blogs/storage/enabling-multiprotocol-workloads-with-amazon-fsx-for-netapp-ontap/

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #935 *Topic 1*

A software company needs to upgrade a critical web application. The application currently runs on a single Amazon EC2 instance that the company hosts in a public subnet. The EC2 instance runs a MySQL database. The application's DNS records are published in an Amazon Route 53 zone.

A solutions architect must reconfigure the application to be scalable and highly available. The solutions architect must also reduce MySQL read latency.

Which combination of solutions will meet these requirements? (Choose two.)

A. Launch a second EC2 instance in a second AWS Region. Use a Route 53 failover routing policy to redirect the traffic to the second EC2 instance.

B. Create and configure an Auto Scaling group to launch private EC2 instances in multiple Availability Zones. Add the instances to a target group behind a new Application Load Balancer. **Most Voted**

C. Migrate the database to an Amazon Aurora MySQL cluster. Create the primary DB instance and reader DB instance in separate Availability Zones. **Most Voted**

D. Create and configure an Auto Scaling group to launch private EC2 instances in multiple AWS Regions. Add the instances to a target group behind a new Application Load Balancer.

E. Migrate the database to an Amazon Aurora MySQL cluster with cross-Region read replicas.

Correct Answer: *BC*

Community vote distribution

BC (95%) DE (5%)

Comments

example_ Highly Voted 10 months, 1 week ago

Selected Answer: BC

https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Concepts.AuroraHighAvailability.html

upvoted 6 times

Sarayounisaldossary Most Recent 2 months ago

Selected Answer: BC

B. Auto Scaling + Load Balancer across AZs
Launch multiple EC2s across multiple AZs

Add to Application Load Balancer

Gives:

HA and Elastic scalability

C. Migrate to Aurora MySQL with a reader in a separate AZ
Aurora is high performance + replication built-in

Reader in another AZ → lower read latency

Built for fault tolerance

NOT A,D,E
A. EC2 in another Region + Route 53 Failover: Disaster recovery, not real-time scaling or latency reduction.

D. Auto Scaling across Regions: Not supported.

E. Aurora with cross-Region replica: Overkill and not needed here.

upvoted 1 times

progounick 9 months, 2 weeks ago

Selected Answer: BC

It is obvious

upvoted 2 times

dhewa 9 months, 3 weeks ago

Selected Answer: BC

These simple options will help you achieve a robust, scalable, and highly available architecture for your web application

upvoted 3 times

[Removed] 9 months, 3 weeks ago

Selected Answer: BC

B and C for high availability and scalability

upvoted 2 times

nebajp 10 months ago

Selected Answer: BC

correct Answer is B& C as it talks about high availability and scalability
and rest of the options are for Disaster Recovery.
Right answer is B & C.

upvoted 3 times

nebajp 10 months ago

correct Answer is B& C as it talks about high availability and scalability
and rest of the options are for Disaster Recovery.
Right answer is B & C.

upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: DE

Answer is DE

upvoted 1 times

JunsK1e 10 months, 1 week ago

Selected Answer: BC

B is correct and C, D is incorrect since the auto scaling is working within Region not Multi region. E incorrect also because the question is asking a high availability C is the best answer that E.

upvoted 4 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #936 *Topic 1*

A company runs thousands of AWS Lambda functions. The company needs a solution to securely store sensitive information that all the Lambda functions use. The solution must also manage the automatic rotation of the sensitive information.

Which combination of steps will meet these requirements with the LEAST operational overhead? (Choose two.)

A. Create HTTP security headers by using Lambda@Edge to retrieve and create sensitive information

B. Create a Lambda layer that retrieves sensitive information **Most Voted**

C. Store sensitive information in AWS Secrets Manager **Most Voted**

D. Store sensitive information in AWS Systems Manager Parameter Store

E. Create a Lambda consumer with dedicated throughput to retrieve sensitive information and create environmental variables

Correct Answer: *BC*

Community vote distribution

BC (90%) CD (10%)

Comments

[Removed] Highly Voted 9 months, 3 weeks ago

Selected Answer: BC

C. Store sensitive information in AWS Secrets Manager.

AWS Secrets Manager securely stores sensitive information and provides automatic rotation of secrets, reducing the need for manual management.
B. Create a Lambda layer that retrieves sensitive information.

Using a Lambda layer allows multiple Lambda functions to access the sensitive information stored in Secrets Manager without needing to duplicate retrieval logic in each function. This approach centralizes the retrieval process and reduces operational complexity.

upvoted 9 times

awisaw 9 months ago

chatgpt answer is now C and D

upvoted 1 times

sOI852POL 7 months, 3 weeks ago

AWS public documentation and other professional forums instead, kindly!
upvoted 3 times

pujithacg8 Highly Voted 10 months ago

D doesn't provide automatic rotation

Answer will be B and C
upvoted 8 times

c3c7e62 Most Recent 1 month ago

Selected Answer: BC

A - not relate, as question does not mention the use of HTTP.
D - technical works, but does not support automatic rotation.
E - not secure to use the environmental variables.

Question and Concern - for the exam, have to pick option B & C. However, really wonder if the approach is practical in real world for secure solution.
1) Secrets Manager encrypt the sensitive info. Lambda Layer saves those sensitive information into /opt layer (in zip format etc) and shares with all other. However, I cannot see there is description the information is encrypted for certain format.
2) After Secret Manager rotation, how can those saved content in /opt folder be refreshed ?

Open for discussion. Thanks.
upvoted 1 times

c3c7e62 1 month ago

https://docs.aws.amazon.com/lambda/latest/dg/chapter-layers.html
upvoted 1 times

Sarayounisaldossary 2 months ago

Selected Answer: BC

A. Lambda@Edge: Not relevant for secrets; Used for CDN/CloudFront.

D. Parameter Store: Doesn't support auto-rotation as easily as Secrets Manager.

E. Env Variables: Not secure and needs manual management.
upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: BC

A - Lambda@Edge is primarily used for managing HTTP request/response behavior in a CDN environment with Amazon CloudFront. It is not designed for securely storing or rotating sensitive information.
B - By creating a layer that securely retrieves sensitive information (for example, from AWS Secrets Manager), we will reduce duplication and operational overhead. The layer can also handle the logic for retrieving updated information if needed.
C - Automatic rotation is managed.
D - Automatic rotation is NOT managed.
E - Setting environment variables dynamically at runtime is not an efficient or scalable practice for "thousands of Lambda functions". Too much operational overhead.
upvoted 5 times

agbor_tambe 8 months, 1 week ago

Selected Answer: BC

BC is correct
upvoted 2 times

progounick 9 months, 1 week ago

Selected Answer: BC

B,C ChatGPT agrees with me
upvoted 1 times

komorebi 10 months, 1 week ago

Selected Answer: CD

Answer is CD

upvoted 1 times

example_ 10 months, 1 week ago

Selected Answer: CD

https://docs.aws.amazon.com/systems-manager/latest/userguide/ps-integration-lambda-extensions.html

https://www.trustwave.com/en-us/resources/blogs/spiderlabs-blog/using-aws-secrets-manager-and-lambda-function-to-store-rotate-and-secure-keys/

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #937 *Topic 1*

A company has an internal application that runs on Amazon EC2 instances in an Auto Scaling group. The EC2 instances are compute optimized and use Amazon Elastic Block Store (Amazon EBS) volumes.

The company wants to identify cost optimizations across the EC2 instances, the Auto Scaling group, and the EBS volumes.

Which solution will meet these requirements with the MOST operational efficiency?

A. Create a new AWS Cost and Usage Report. Search the report for cost recommendations for the EC2 instances the Auto Scaling group, and the EBS volumes.

B. Create new Amazon CloudWatch billing alerts. Check the alert statuses for cost recommendations for the EC2 instances, the Auto Scaling group, and the EBS volumes.

C. Configure AWS Compute Optimizer for cost recommendations for the EC2 instances, the Auto Scaling group and the EBS volumes. **Most Voted**

D. Configure AWS Compute Optimizer for cost recommendations for the EC2 instances. Create a new AWS Cost and Usage Report. Search the report for cost recommendations for the Auto Scaling group and the EBS volumes.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

Sarayounisaldossary 2 months ago

Selected Answer: C

the question here focuses on (COST Optimization)

- AWS Cost and Usage Report is the most detailed report that shows your AWS usage and charges.

- cloudwatch is only for setting alerts based on setting policies-

- would go right-away with AWS compute Optimizer because this service gives recommendations to optimize your compute resources (like EC2, Lambda, etc.) for better performance and lower cost.

which is C and D

BUT when reading D choice there's nothing wrong with AWS Cost and Usage report but why would i make it more complex to get the cost optimization? AWS compute Optimizer is more than enough which is C option

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: C

Go with Option C for the most efficient and straightforward solution for sure. Just one more thing: AWS Compute Optimizer doesn’t provide cost data directly (such as detailed dollar amounts), so supplementing it with AWS Cost Explorer or Trusted Advisor is a good practice for deeper cost analysis.

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: C

C looks right

upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: C

Answer is C

upvoted 2 times

example_ 10 months, 1 week ago

Selected Answer: C

https://aws.amazon.com/compute-optimizer/
https://docs.aws.amazon.com/compute-optimizer/latest/ug/what-is-compute-optimizer.html

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #938 *Topic 1*

A company is running a media store across multiple Amazon EC2 instances distributed across multiple Availability Zones in a single VPC. The company wants a high-performing solution to share data between all the EC2 instances, and prefers to keep the data within the VPC only.

What should a solutions architect recommend?

A. Create an Amazon S3 bucket and call the service APIs from each instance's application

B. Create an Amazon S3 bucket and configure all instances to access it as a mounted volume

C. Configure an Amazon Elastic Block Store (Amazon EBS) volume and mount it across all instances

D. Configure an Amazon Elastic File System (Amazon EFS) file system and mount it across all instances **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

Sarayounisaldossary 2 months ago

Selected Answer: D

Why is EFS the correct answer?

Amazon EFS is a fully managed, shared file system.

It can be mounted on multiple EC2 instances at the same time.

It works across multiple Availability Zones.

High performance, and keeps all data within the VPC.

Perfect fit for the use case: “Share data between EC2 instances.”

Other options:

S3 is object storage, not a shared file system.

EBS can be attached to only one instance at a time.

Mounting S3 isn’t officially supported or reliable.

upvoted 1 times

aragon_saa 9 months, 4 weeks ago

Selected Answer: D

Answer is D

upvoted 4 times

muhammadahmer36 9 months, 4 weeks ago

Selected Answer: D

D is the right answer

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #939 *Topic 1*

A company uses an Amazon RDS for MySQL instance. To prepare for end-of-year processing, the company added a read replica to accommodate extra read-only queries from the company's reporting tool. The read replica CPU usage was 60% and the primary instance CPU usage was 60%.

After end-of-year activities are complete, the read replica has a constant 25% CPU usage. The primary instance still has a constant 60% CPU usage. The company wants to rightsize the database and still provide enough performance for future growth.

Which solution will meet these requirements?

A. Delete the read replica Do not make changes to the primary instance

B. Resize the read replica to a smaller instance size Do not make changes to the primary instance

C. Resize the read replica to a larger instance size Resize the primary instance to a smaller instance size

D. Delete the read replica Resize the primary instance to a larger instance **Most Voted**

Correct Answer: *D*

Community vote distribution

D (54%) B (46%)

Comments

Oghare Highly Voted 8 months ago

Answer is B
Since the read replica is now underutilized with only 25% CPU usage, it can be resized to a smaller instance to save costs while still handling the reduced read queries. No changes to the primary instance are needed, as it is consistently running at 60% CPU usage, which is manageable, and the read replica will still offload read queries in the future.

upvoted 7 times

7dcef09 Most Recent 1 month, 3 weeks ago

Selected Answer: D

Delete the read replica and resize the primary instance seems to be the right choice for this question

upvoted 1 times

Saravounisaldossary 2 months ago

Selected Answer: D

You should remove an underused resource (read replica).

Use that budget to increase primary capacity, solving the CPU load.

Even though Option B is “safe” and cost-conscious, it ignores the performance needs of the primary, which is where the real issue lies.
Option D is more aligned with AWS best practices and the question's goal of performance + future growth.

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: B

Not A or D because they risk future scalability issues during traffic spikes. Moreover, when the end of next year comes, you have to bring back the deleted instance and configure it all over again. As a solutions architect, I wouldn't recommend things like that (personally).

upvoted 2 times

Ghoneam 5 months, 3 weeks ago

Selected Answer: D

D. Delete the read replica Resize the primary instance to a larger instance
because
no need for read replica after year end
still provide enough performance for future growth.

upvoted 4 times

Abdullah2004 9 months, 1 week ago

No clear explanation here and everyone agreed with default answer

upvoted 1 times

[Removed] 9 months, 3 weeks ago

Selected Answer: B

B looks good

upvoted 1 times

aragon_saa 9 months, 4 weeks ago

Selected Answer: B

Answer is B

upvoted 1 times

muhammadahmer36 9 months, 4 weeks ago

Selected Answer: B

Resize the read replica to a smaller instance size Do not make changes to the primary instance

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #940 *Topic 1*

A company is migrating its databases to Amazon RDS for PostgreSQL. The company is migrating its applications to Amazon EC2 instances. The company wants to optimize costs for long-running workloads.

Which solution will meet this requirement MOST cost-effectively?

A. Use On-Demand Instances for the Amazon RDS for PostgreSQL workloads. Purchase a 1 year Compute Savings Plan with the No Upfront option for the EC2 instances.

B. Purchase Reserved Instances for a 1 year term with the No Upfront option for the Amazon RDS for PostgreSQL workloads. Purchase a 1 year EC2 Instance Savings Plan with the No Upfront option for the EC2 instances.

C. Purchase Reserved Instances for a 1 year term with the Partial Upfront option for the Amazon RDS for PostgreSQL workloads. Purchase a 1 year EC2 Instance Savings Plan with the Partial Upfront option for the EC2 instances.

D. Purchase Reserved Instances for a 3 year term with the All Upfront option for the Amazon RDS for PostgreSQL workloads. Purchase a 3 year EC2 Instance Savings Plan with the All Upfront option for the EC2 instances. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

LeonSauveterre Highly Voted 5 months ago

Selected Answer: D

What you should know:
1. No Upfront - Pay nothing at the time of purchase. Payments are spread out monthly over the term (1 or 3 years). Lowest savings with high flexibility.
2. Partial Upfront - Pay a portion of the cost upfront, with reduced monthly payments over the term. Moderate savings with moderate flexibility.
3. All Upfront - Pay the full cost upfront at the time of purchase. No additional payments required during the term. Highest savings with the least flexibility.

Merely by the question, option D would be the answer because it's best for companies with a stable workload, long-term commitment, and sufficient budget for upfront payments.

Just to add something, generally option C would be the best in real world cases because it strikes a good balance between savings and upfront investment, and uitable for mid-term commitments with moderate budget flexibility.

upvoted 6 times

LeonSauveterre 5 months ago

uitable => suitable sorry

upvoted 1 times

744fdad Most Recent 9 months, 4 weeks ago

the longer the reservation the more the savings too for reserved

upvoted 2 times

muhammadahmer36 9 months, 4 weeks ago

Selected Answer: D

D is right

upvoted 1 times

nebajp 10 months ago

Selected Answer: D

Correct Answer - D,
All upfront is cheaper than partial and no upfront.

upvoted 1 times

swati1508 10 months ago

D is correct

upvoted 2 times

JunsK1e 10 months, 1 week ago

Selected Answer: D

letter D is the correct, It gives you big discount if you purchase with all upfront

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #941 *Topic 1*

A company is using an Amazon Elastic Kubernetes Service (Amazon EKS) cluster. The company must ensure that Kubernetes service accounts in the EKS cluster have secure and granular access to specific AWS resources by using IAM roles for service accounts (IRSA).

Which combination of solutions will meet these requirements? (Choose two.)

A. Create an IAM policy that defines the required permissions Attach the policy directly to the IAM role of the EKS nodes.

B. Implement network policies within the EKS cluster to prevent Kubernetes service accounts from accessing specific AWS services.

C. Modify the EKS cluster's IAM role to include permissions for each Kubernetes service account. Ensure a one-to-one mapping between IAM roles and Kubernetes roles.

D. Define an IAM role that includes the necessary permissions. Annotate the Kubernetes service accounts with the Amazon ResourceName (ARN) of the IAM role. **Most Voted**

E. Set up a trust relationship between the IAM roles for the service accounts and an OpenID Connect (OIDC) identity provider. **Most Voted**

Correct Answer: *DE*

Community vote distribution

DE (100%)

Comments

Dantecito 3 months ago

Selected Answer: DE

This was in my Exam today. Good Luck!!!!!!!!!!!!

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: DE

A - Attaching policies to the IAM role of the EKS nodes provides broad permissions to all pods running on the nodes. IRSA is designed to eliminate such over-permissive configurations.
B - AWS resource access is managed by IAM, not Kubernetes network policies.

C - EKS cluster's IAM role is not meant to handle granular access for service accounts.

D - Ensures fine-grained access control at the service account level.
E - Allows Kubernetes service accounts to assume IAM roles securely.

upvoted 2 times

spoved 8 months, 2 weeks ago

https://docs.aws.amazon.com/eks/latest/userguide/iam-roles-for-service-accounts.html
https://docs.aws.amazon.com/eks/latest/userguide/associate-service-account-role.html
https://docs.aws.amazon.com/eks/latest/userguide/enable-iam-roles-for-service-accounts.html
=> DE

upvoted 4 times

mooondooo 8 months, 2 weeks ago

Selected Answer: DE

probably D and E

https://docs.aws.amazon.com/emr/latest/EMR-on-EKS-DevelopmentGuide/setting-up-enable-IAM.html

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: DE

chatgpt: D. Define an IAM role that includes the necessary permissions. Annotate the Kubernetes service accounts with the Amazon Resource Name (ARN) of the IAM role:

Granular Access Control: By defining an IAM role with the necessary permissions and annotating the Kubernetes service accounts with the ARN of this IAM role, you can achieve fine-grained access control for specific AWS resources. This allows each service account to have only the permissions it needs.
E. Set up a trust relationship between the IAM roles for the service accounts and an OpenID Connect (OIDC) identity provider:

IRSA Integration: To enable IRSA, your EKS cluster must be associated with an OpenID Connect (OIDC) identity provider. This trust relationship allows Kubernetes service accounts to assume IAM roles, enabling secure and granular access to AWS resources.

upvoted 2 times

744fdad 9 months, 3 weeks ago

my educated guess C, E

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #942 *Topic 1*

A company regularly uploads confidential data to Amazon S3 buckets for analysis.

The company's security policies mandate that the objects must be encrypted at rest. The company must automatically rotate the encryption key every year. The company must be able to track key rotation by using AWS CloudTrail. The company also must minimize costs for the encryption key.

Which solution will meet these requirements?

A. Use server-side encryption with customer-provided keys (SSE-C)

B. Use server-side encryption with Amazon S3 managed keys (SSE-S3)

C. Use server-side encryption with AWS KMS keys (SSE-KMS) **Most Voted**

D. Use server-side encryption with customer managed AWS KMS keys

Correct Answer: *C*

Community vote distribution

C (64%) D (36%)

Comments

nebajp Highly Voted 10 months ago

Selected Answer: C

SSE keys provided usage fee application and there is no monthly charges, hence its a correct option.
D is highly cost option with monthly and usage fee. which is incorrect.

upvoted 10 times

tch Most Recent 3 months ago

Selected Answer: C

should put full name as "Use server-side encryption with AWS KMS managed keys"

answer D does not minimize the cost

upvoted 1 times

yangbo 3 months ago

Selected Answer: D

"SSE-KMS (AWS managed KMS keys): Automatically rotated, but the rotation records are not visible. Users cannot view the RotateKey event in CloudTrail."

upvoted 2 times

Dantecito 3 months, 2 weeks ago

Selected Answer: D

C. since May 2022 managed SSE-KMS can rotate every year https://docs.aws.amazon.com/kms/latest/developerguide/rotate-keys.html
D. I still prefer D because the question says "The company must automatically rotate the encryption key every year" it does not say "AWS must automatically..."

upvoted 2 times

tch 3 months ago

customer managed keys has monthly fee for existence of keys (pro-rated hourly). Also charged for key usage.... not an good idea ....

upvoted 1 times

GOTJ 4 months ago

Selected Answer: C

With all due respect, I think option "C" is poorly written, and it reflects on the lack of consensus with this one. A "server-side encryption with customer managed AWS KMS keys" (option "D") is nothing but an implementation of a "server-side encryption with AWS KMS keys" (option "C"), and both are SSE-KMS. What I think op tried to express with option C is: "Use server-side encryption with AWS-MANAGED AWS KMS keys (SSE-KMS). If my guess is correct, option "C" is the right answer, according to the table shown in this link: https://docs.aws.amazon.com/kms/latest/developerguide/concepts.html#aws-managed-cmk

Option C and Option D satisfied all the technical requirements of the company, but option C is $1/month cheaper than option D

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: C

AWS-managed keys are automatically rotated every year (approximately 365 days) by AWS.
Key rotation cannot be enabled or disabled for AWS-managed keys, ensuring compliance with the rotation requirement.

When AWS KMS rotates the key material for an AWS-managed key, it writes: A KMS CMK Rotation event to Amazon EventBridge, A RotateKey event to AWS CloudTrail.
This allows the company to track key rotation events for auditing purposes.

Cost Minimization:AWS-managed keys are free to use (no additional cost beyond the base AWS KMS pricing for API calls). This helps minimize costs.
https://docs.aws.amazon.com/kms/latest/developerguide/rotate-keys.html

upvoted 2 times

FlyingHawk 4 months, 2 weeks ago

Why not D?
While customer-managed KMS keys also support automatic rotation and CloudTrail logging, they incur additional costs:
Customer-managed keys have a monthly fee (e.g., $1 per key per month).
This would increase costs compared to AWS-managed keys, which are free.
Since the requirement is to minimize costs, AWS-managed keys are the better choice.

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

You can enable automatic key rotation on the customer managed keys that you use for server-side encryption in AWS services.
if disabled, then won't meet the requirement "The company must automatically rotate the encryption key every year"

upvoted 1 times

joanna91 4 months, 3 weeks ago

Selected Answer: C

"Automatically" rotate, then make it handled by AWS service, not by customer -> C, not D

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: D

"automatically rotate the encryption key" => C or D (because of KMS), then
"able to track key rotation" => Just D.

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Monitoring key rotation
When AWS KMS rotates the key material for an AWS managed key or customer managed key, it writes a KMS CMK Rotation event to Amazon EventBridge and a RotateKey event to your AWS CloudTrail log. You can use these records to verify that the KMS key was rotated.

upvoted 1 times

Anyio 5 months, 2 weeks ago

Selected Answer: D

Option C: Incorrect. Though server-side encryption with AWS KMS keys (SSE-KMS) would allow AWS to manage keys and enable logging via AWS CloudTrail, this option uses AWS-managed keys instead of customer-managed keys, limiting control over key rotations. Additionally, there can be more costs involved in using AWS-managed KMS keys compared to the customer managing their own.

Option D: Correct. Using server-side encryption with customer-managed AWS KMS keys allows the company to have full control over the encryption keys, including managing and ensuring automatic rotation every year. Moreover, AWS CloudTrail can be employed to log events associated with AWS KMS, enabling the tracking of when keys are rotated. This option balances cost-effectiveness with the operational requirements specified, as it provides the necessary control without unnecessary expenses from more specialized AWS services.

upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: D

I will choose Option D for the following reasons:

#1 Automatic key rotation: AWS KMS allows you to set up automatic key rotation for customer managed keys, which fulfills the requirement to rotate encryption keys yearly.

2 CloudTrail tracking: All KMS key operations are logged in CloudTrail, enabling you to track key rotation activity.

#3 Lowest cost: While using customer-provided keys (SSE-C) might seem cost-effective at first glance, managing your own keys adds complexity and can be more expensive in the long run.

#$ Compliance with security policies: Using customer managed KMS keys ensures that the company has full control over the encryption keys, meeting the stringent security requirements

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

There is no requirement for the company has full control over the encryption keys. the requirement is must auto-rotate the keys every year, must be able to track rotation via cloudtrail, and minimise cost. C meets requirement as AWS KMS with AWS Managed Key auto-rotate yearly too, you cannot change the rotation period or disable the rotation, it also logs the rotation events and no cost for storing it. you only need the customer key if you want the flexibility to delete it or change rotation period or disable it... etc but this is not required in the problem.

upvoted 1 times

XXXXXlNN 7 months, 3 weeks ago

D
auto-rotation feature > customer managed key

upvoted 1 times

XXXXXlNN 8 months, 1 week ago

D. customer needs to see the logs from Cloudtrail!

upvoted 1 times

sOI852POL 7 months, 3 weeks ago

Even with AWS KMS keys, rotation is logged on ctrail. Answer is D.

https://docs.aws.amazon.com/kms/latest/developerguide/rotate-keys.html#:~:text=Monitoring%20key%20rotation,key%20was%20rotated.

upvoted 1 times

sOI852POL 9 months ago

Selected Answer: C

Answer is C.
There is no monthly fee for AWS managed keys
https://docs.aws.amazon.com/kms/latest/developerguide/concepts.html#aws-managed-cmk

upvoted 3 times

elmyth 9 months, 1 week ago

Selected Answer: C

Customer managed key: Monthly fee (pro-rated hourly) + Per-use fee + rotation and cloudtrail
AWS managed key: No monthly fee + Per-use fee (some AWS services pay this fee for you)+ rotation and cloudtrail

upvoted 4 times

dhewa 9 months, 3 weeks ago

Selected Answer: D

D gives you control, allows you to customise for example rotation policies to suit your compliance needs.

upvoted 2 times

komorebi 10 months ago

Selected Answer: D

Answer is D

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #943 *Topic 1*

A company has migrated several applications to AWS in the past 3 months. The company wants to know the breakdown of costs for each of these applications. The company wants to receive a regular report that includes this information.

Which solution will meet these requirements MOST cost-effectively?

A. Use AWS Budgets to download data for the past 3 months into a .csv file. Look up the desired information.

B. Load AWS Cost and Usage Reports into an Amazon RDS DB instance. Run SQL queries to get the desired information.

C. Tag all the AWS resources with a key for cost and a value of the application's name. Activate cost allocation tags. Use Cost Explorerto get the desired information. **Most Voted**

D. Tag all the AWS resources with a key for cost and a value of the application's name. Use the AWS Billing and Cost Management console todownload bills for the past 3 months. Look up the desired information.

Correct Answer: *C*

Community vote distribution

C (80%) D (20%)

Comments

spoved 8 months, 1 week ago

Selected Answer: C

- Organizing and tracking costs using AWS cost allocation tags
https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/cost-alloc-tags.html
- Cost Explorer uses the same dataset that is used to generate the AWS Cost and Usage Reports and the detailed billing reports. For a comprehensive review of the data, you can download it into a comma-separated value (CSV) file.
https://docs.aws.amazon.com/cost-management/latest/userguide/ce-what-is.html

upvoted 4 times

rpmaws 8 months, 3 weeks ago

Selected Answer: C

cost explorer allows cost analysis of up to 13 months.

upvoted 2 times

ccceb01 9 months, 1 week ago

Selected Answer: D

https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/cost-alloc-tags.html
upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: C

c looks good
upvoted 2 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #944 *Topic 1*

An ecommerce company is preparing to deploy a web application on AWS to ensure continuous service for customers. The architecture includes a web application that the company hosts on Amazon EC2 instances, a relational database in Amazon RDS, and static assets that the company stores in Amazon S3.

The company wants to design a robust and resilient architecture for the application.

Which solution will meet these requirements?

A. Deploy Amazon EC2 instances in a single Availability Zone. Deploy an RDS DB instance in the same Availability Zone. Use Amazon S3 with versioning enabled to store static assets.

B. Deploy Amazon EC2 instances in an Auto Scaling group across multiple Availability Zones. Deploy a Multi-AZ RDS DB instance. Use Amazon CloudFront to distribute static assets. **Most Voted**

C. Deploy Amazon EC2 instances in a single Availability Zone. Deploy an RDS DB instance in a second Availability Zone for cross-AZ redundancy. Serve static assets directly from the EC2 instances.

D. Use AWS Lambda functions to serve the web application. Use Amazon Aurora Serverless v2 for the database. Store static assets in Amazon Elastic File System (Amazon EFS) One Zone-Infrequent Access (One Zone-IA).

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

mk168898 6 months, 4 weeks ago

keyword is multi-AZ

upvoted 2 times

Jeyaluxshan 9 months ago

B is the answer

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: B

B is the answer

upvoted 3 times

pujithacg8 10 months ago

B is correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #945 *Topic 1*

An ecommerce company runs several internal applications in multiple AWS accounts. The company uses AWS Organizations to manage its AWS accounts.

A security appliance in the company's networking account must inspect interactions between applications across AWS accounts.

Which solution will meet these requirements?

A. Deploy a Network Load Balancer (NLB) in the networking account to send traffic to the security appliance. Configure the application accounts to send traffic to the NLB by using an interface VPC endpoint in the application accounts.

B. Deploy an Application Load Balancer (ALB) in the application accounts to send traffic directly to the security appliance.

C. Deploy a Gateway Load Balancer (GWLB) in the networking account to send traffic to the security appliance. Configure the application accounts to send traffic to the GWLB by using an interface GWLB endpoint in the application accounts.
Most Voted

D. Deploy an interface VPC endpoint in the application accounts to send traffic directly to the security appliance.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

mk168898 Highly Voted 6 months, 4 weeks ago

security appliance, inspect interaction => GWLB

upvoted 5 times

poramesh Most Recent 3 months, 1 week ago

Finallly a ques on GLB

upvoted 3 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: C

https://docs.aws.amazon.com/elasticloadbalancing/latest/gateway/getting-started.html
upvoted 1 times

[Removed] 9 months, 3 weeks ago

Selected Answer: C

C sounds right
upvoted 2 times

swati1508 10 months ago

Inspect traffic for that use GWLB OPTION C
upvoted 3 times

komorebi 10 months, 1 week ago

Selected Answer: C

Answer is C
upvoted 2 times

example_ 10 months, 1 week ago

Selected Answer: C

https://docs.aws.amazon.com/vpc/latest/privatelink/create-gateway-load-balancer-endpoint-service.html
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #946 *Topic 1*

A company runs its production workload on an Amazon Aurora MySQL DB cluster that includes six Aurora Replicas. The company wants near-real-time reporting queries from one of its departments to be automatically distributed across three of the Aurora Replicas. Those three replicas have a different compute and memory specification from the rest of the DB cluster.

Which solution meets these requirements?

A. Create and use a custom endpoint for the workload **Most Voted**

B. Create a three-node cluster clone and use the reader endpoint

C. Use any of the instance endpoints for the selected three nodes

D. Use the reader endpoint to automatically distribute the read-only workload

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

[Removed] Highly Voted 9 months, 3 weeks ago

Selected Answer: A

Custom Endpoints:

Custom endpoints in Amazon Aurora allow you to group specific replicas together and route traffic only to those replicas. This is particularly useful when you have replicas with different compute and memory specifications and want to direct specific workloads, such as reporting queries, to those replicas.
By creating a custom endpoint, you can include the three specific Aurora Replicas that have the required compute and memory configurations, ensuring that your near-real-time reporting queries are automatically distributed among these replicas.

upvoted 7 times

FlyingHawk Most Recent 4 months, 2 weeks ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.Endpoints.Custom.html
A custom endpoint for an Aurora cluster represents a set of DB instances that you choose. When you connect to the endpoint, Aurora performs connection balancing and chooses one of the instances in the group to handle the connection. You define which instances this endpoint refers to, and you decide what purpose the endpoint serves.

upvoted 2 times