

Configuration changes= AWS Config

upvoted 34 times

ellanos Most Recent 1 month, 3 weeks ago

Selected Answer: A

A is the most suitable option

upvoted 1 times

Juju66 1 month, 3 weeks ago

Selected Answer: A

c - AWS Inspector is only for EC2 instances

upvoted 2 times

satyaamm 5 months, 2 weeks ago

Selected Answer: A

AWS Config is the only one that provides continuous monitoring here.

upvoted 1 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: A

To review an AWS deployment and ensure that Amazon S3 buckets do not have unauthorized configuration changes, you need a service that can monitor and record the configuration state of AWS resources and evaluate them against desired configurations.

upvoted 1 times

PaulGa 10 months ago

Selected Answer: A

Ans A - as well explained by "Buruguduystunstugudunstuy" – we are dealing with configuration here: ensuring that what we've designed continues to follow the rules

upvoted 3 times

andyngkh86 1 year, 5 months ago

ChatGPT give the answer is D

upvoted 2 times

A_jaa 1 year, 5 months ago

Selected Answer: A

Answer-A

upvoted 1 times

Ruffyit 1 year, 7 months ago

A: <https://aws.amazon.com/config/#:~:text=How%20it%20works-,AWS%20Config,-continually%20assesses%2C%20audits>

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: A

AWS Config continually assesses, audits, and evaluates the configurations and relationships of your resources on AWS, on premises, and on other clouds. It normalizes changes into a consistent format and checks resource compliance with custom and managed rules before and after provisioning.

<https://aws.amazon.com/config/#:~:text=How%20it%20works-,AWS%20Config,-continually%20assesses%2C%20audits>

upvoted 4 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: A

AWS Config provides a detailed inventory of the company's AWS resources and configuration history, and can be configured with rules to evaluate resource configurations for compliance with policies and best practices.

The solutions architect can enable AWS Config and configure rules specifically checking for S3 bucket settings like public access blocking, encryption settings, access control lists, etc. AWS Config will record configuration changes to S3 buckets over time, allowing the company to review changes and be alerted about any unauthorized modifications.

By. Claude.ai

upvoted 2 times

miki111 1 year, 11 months ago

Option A is the right answer for this.

upvoted 1 times

cookieMr 1 year, 12 months ago

Selected Answer: A

AWS Config is a service that provides a detailed view of the configuration of AWS resources in your account. By enabling AWS Config, you can capture configuration changes and maintain a record of resource configurations over time. It allows you to define rules that check for compliance with desired configurations and can generate alerts or automated actions when unauthorized changes occur.

To accomplish the goal of preventing unauthorized configuration changes in Amazon S3 buckets, you can configure AWS Config rules specifically for S3 bucket configurations. These rules can check for a variety of conditions, such as ensuring that encryption is enabled, access control policies are correctly configured, and public access is restricted.

While options B, C, and D offer valuable services for various aspects of AWS deployment, they are not specifically focused on preventing unauthorized configuration changes in Amazon S3 buckets as effectively as enabling AWS Config.

upvoted 2 times

Abrar2022 2 years, 1 month ago

Don't be mistaken in thinking that it's Server access logs because that's for detailed records for requests made to S3. It's AWS Config because it records configuration changes.

upvoted 1 times

Rahulbit34 2 years, 1 month ago

AWS Trusted Advisor is for providing recommendation only.

For any configuration use AWS Config

Inspector is for scanning for any software vulnerabilities and unintended network exposure

upvoted 1 times

PhucVuu 2 years, 1 month ago

Selected Answer: A

To accomplish the goal of ensuring that Amazon S3 buckets do not have unauthorized configuration changes, a solutions architect should turn on AWS Config with the appropriate rules. AWS Config enables continuous monitoring and recording of AWS resource configurations, including S3 buckets. By turning on AWS Config with the appropriate rules, the solutions architect can be notified of any unauthorized changes made to the S3 bucket configurations, allowing for prompt corrective action. Options B, C, and D are not directly related to monitoring and preventing unauthorized configuration changes to Amazon S3 buckets.

upvoted 1 times


Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #27

Topic 1

A company is launching a new application and will display application metrics on an Amazon CloudWatch dashboard. The company's product manager needs to access this dashboard periodically. The product manager does not have an AWS account. A solutions architect must provide access to the product manager by following the principle of least privilege. Which solution will meet these requirements?

- A. Share the dashboard from the CloudWatch console. Enter the product manager's email address, and complete the sharing steps. Provide a shareable link for the dashboard to the product manager. **Most Voted**
- B. Create an IAM user specifically for the product manager. Attach the CloudWatchReadOnlyAccess AWS managed policy to the user. Share the new login credentials with the product manager. Share the browser URL of the correct dashboard with the product manager.
- C. Create an IAM user for the company's employees. Attach the ViewOnlyAccess AWS managed policy to the IAM user. Share the new login credentials with the product manager. Ask the product manager to navigate to the CloudWatch console and locate the dashboard by name in the Dashboards section.
- D. Deploy a bastion server in a public subnet. When the product manager requires access to the dashboard, start the server and share the RDP credentials. On the bastion server, ensure that the browser is configured to open the dashboard URL with cached AWS credentials that have appropriate permissions to view the dashboard.

Correct Answer: A*Community vote distribution***Comments****masetromain** **Highly Voted** 2 years, 8 months ago**Selected Answer: A**Answer A : <https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/cloudwatch-dashboard-sharing.html>

Share a single dashboard and designate specific email addresses of the people who can view the dashboard. Each of these users creates their own password that they must enter to view the dashboard.

upvoted 95 times

123jhl0 2 years, 8 months ago

Thanks for the help. No doubt A is the answer.

I thanks for the link! NO doubt A is the answer.

upvoted 10 times

mn2013 1 year, 4 months ago

But this link also says All people who you share the dashboard with are granted these permissions for the account. If you share the dashboard publicly, then everyone who has the link to the dashboard has these permissions.

The cloudwatch:GetMetricData and ec2:DescribeTags permissions cannot be scoped down to specific metrics or EC2 instances, so the people with access to the dashboard can query all CloudWatch metrics and the names and tags of all EC2 instances in the account. If that is the case, how is the least privilege principle applicable?

upvoted 4 times

Yak_Yeti 2 months, 2 weeks ago

while CloudWatch dashboard sharing doesn't perfectly align with the strictest interpretation of least privilege, it can be made more secure by carefully designing the dashboard, implementing strong governance, and educating the user. If the data within the dashboard is highly sensitive, then creating a custom application would be the best method.

Therefore, while option A is the most efficient and simple, it is important to understand the security trade offs that are being made.

upvoted 1 times

Guru4Cloud **Highly Voted** 1 year, 10 months ago

Selected Answer: B

Option B provides the product manager with specific access to the CloudWatch dashboard using an IAM user with the CloudWatchReadOnlyAccess policy attached. The IAM user has only read-only access to the required resources, which follows the principle of least privilege.

upvoted 19 times

emilyhu08 1 year, 8 months ago

b has a problem for cloudwatchreadonlyaccess policy, it's not only grant read access to dashboard, but other read permission for logs, insights, etc. so it does not follows the principle of least privilege. Option A only grants access to dashboard.

upvoted 18 times

Juju66 **Most Recent** 1 month, 3 weeks ago

Selected Answer: A

This is new function for cloutwatch now.

upvoted 2 times

Vandaman 3 months, 3 weeks ago

Selected Answer: A

Answer is A. This option is available in Cloudwatch

upvoted 1 times

Nacho2222222 4 months ago

Selected Answer: B

CloudWatch does not allow public links by default, but you can share temporary access using AWS IAM Identity Center (formerly AWS SSO) or a federated user with specific permissions.

upvoted 1 times

Krenil 4 months, 1 week ago

Selected Answer: B

Option - B:

Here's why:

Principle of Least Privilege: The product manager should only be given the permissions necessary to access the CloudWatch dashboard. The CloudWatchReadOnlyAccess policy grants the least privilege access to view CloudWatch metrics and dashboards.

IAM User: Since the product manager does not have an AWS account, creating a specific IAM user for the product manager ensures they can log in securely to AWS, without needing to share broader permissions.

Simple Access: After the IAM user is created, the product manager can directly log into AWS, navigate to the CloudWatch dashboard, and access the metrics they need, without excessive complexity or unnecessary access.

upvoted 1 times

testiaced 3 months ago

"To view a dashboard that you create, users must have the required permissions, or the dashboard must be shared across AWS accounts. You can't just email a link to a user who isn't signed into AWS and allow them to view the dashboard."

upvoted 1 times

AfricanCloudGuru 4 months, 2 weeks ago

Selected Answer: A

Since the Product manager access the account periodically , sharing the account will be the best option

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/cloudwatch-dashboard-sharing.html>

upvoted 1 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: A

CloudWatch allows dashboards to be shared publicly or with specific users using a shareable URL. This enables access without requiring an AWS account.

This approach follows the principle of least privilege because the product manager is given only the ability to view the dashboard via the shared link, without needing broader access to the AWS environment.

This is the correct choice because it directly meets the requirements without granting unnecessary permissions or requiring the creation of AWS credentials.

upvoted 3 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: A

1.Direct Sharing Without AWS Account:

Amazon CloudWatch allows dashboards to be shared externally via a shareable link. This enables users without an AWS account to access the dashboard securely, which meets the requirement to allow the product manager (who doesn't have an AWS account) to view the dashboard.

2.Least Privilege Principle:

The product manager only gets access to the specific dashboard, and no additional permissions or access to the AWS environment are granted. This aligns with the principle of least privilege.

3.No Operational Overhead:

The solution does not require creating new IAM users or setting up infrastructure like bastion servers. It's straightforward and efficient.

upvoted 1 times

EzKkk 6 months, 3 weeks ago

Selected Answer: B

I will go with B even though A sounds like a better choice for the following reasons:

- 1, You need a solution to share your dashboard.
- 2, You only share a specific metric of an application to whom lack access to AWS.
- 3, Target in question only access periodically.
- 4, Must apply least privilege.

Though A & B are good choices to share you dashboard, A doesn't follow least privilege principal because it provides more permission than needed. Not to mention this solution also raises concerns about security since this dashboard is public.

upvoted 2 times

f51a8bd 7 months ago

respuesta es A:

Esta solución permite cumplir con los requisitos de acceso seguro y controlado sin una cuenta completa de AWS y sigue el principio de mínimo privilegio para acceder solo a los paneles necesarios en Amazon CloudWatch.

upvoted 1 times

mzeynalli 7 months, 3 weeks ago

Selected Answer: A

The correct answer is:

A. Share the dashboard from the CloudWatch console. Enter the product manager's email address, and complete the sharing steps. Provide a shareable link for the dashboard to the product manager.

Explanation:

Principle of Least Privilege: The solution should only provide the minimum necessary access. Sharing a CloudWatch dashboard directly meets this principle without the need for an AWS account.

Shareable Link: CloudWatch allows you to share a dashboard using a shareable URL that can be accessed without requiring AWS credentials. This way, the product manager can view the dashboard without having an AWS account.

No IAM User Required: Solutions involving IAM users (options B and C) require creating an AWS identity and sharing credentials, which introduces unnecessary security risks and management overhead.

No Bastion Server: Deploying a bastion server (option D) is not practical, involves additional costs, and requires managing another infrastructure component, which goes against simplicity and security best practices.

upvoted 3 times

mzeynalli 7 months, 3 weeks ago

The correct answer is:

A. Share the dashboard from the CloudWatch console. Enter the product manager's email address, and complete the sharing steps. Provide a shareable link for the dashboard to the product manager.

Explanation:

Principle of Least Privilege: The solution should only provide the minimum necessary access. Sharing a CloudWatch dashboard directly meets this principle without the need for an AWS account.

Shareable Link: CloudWatch allows you to share a dashboard using a shareable URL that can be accessed without requiring AWS credentials. This way, the product manager can view the dashboard without having an AWS account.

No IAM User Required: Solutions involving IAM users (options B and C) require creating an AWS identity and sharing credentials, which introduces unnecessary security risks and management overhead.

No Bastion Server: Deploying a bastion server (option D) is not practical, involves additional costs, and requires managing another infrastructure component, which goes against simplicity and security best practices.

upvoted 2 times

e854fbe 8 months ago

Selected Answer: B

correct answer B

the reason A is not a good answer is:

when sharing a link to the CloudWatch dashboard the following warning appears

We recommend that you do not share dashboards if your account contains any sensitive information which you would not wish to share with the users with whom you are sharing the dashboard.

The users that you specified above will be granted the following permissions: CloudWatch read-only permissions to alarms and contributor insights rules in the Dashboard which you share, and to all metrics and the names and tags of all EC2 instances in your account even if they are not shown in the Dashboard which you share. We recommend that you consider whether it is appropriate to make this information available to the users with whom you are sharing.

so following the least privilege principle, creating an IAM User in option B is more secure

upvoted 3 times

Micha3l 8 months, 4 weeks ago

Selected Answer: B

The goal is to allow a product manager (who does not have an AWS account) to access a CloudWatch dashboard periodically.

B follows the principle of least privilege, ensuring that the product manager can only view the dashboard and not perform any other actions within AWS. Is direct email sharing is not a feature of CloudWatch?

upvoted 2 times

Omshanti 9 months ago

Selected Answer: B

Least privilege

upvoted 1 times

pranavff_examtopics_1993 9 months, 1 week ago

Selected Answer: B

A would be indeed following the principle of least privilege, but periodic access means we have to do this time and time again when the product manager requests access. B is better

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #28

Topic 1

A company is migrating applications to AWS. The applications are deployed in different accounts. The company manages the accounts centrally by using AWS Organizations. The company's security team needs a single sign-on (SSO) solution across all the company's accounts. The company must continue managing the users and groups in its on-premises self-managed Microsoft Active Directory.

Which solution will meet these requirements?

- A. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console. Create a one-way forest trust or a one-way domain trust to connect the company's self-managed Microsoft Active Directory with AWS SSO by using AWS Directory Service for Microsoft Active Directory.
- B. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console. Create a two-way forest trust to connect the company's self-managed Microsoft Active Directory with AWS SSO by using AWS Directory Service for Microsoft Active Directory. **Most Voted**
- C. Use AWS Directory Service. Create a two-way trust relationship with the company's self-managed Microsoft Active Directory.
- D. Deploy an identity provider (IdP) on premises. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console.

Correct Answer: B*Community vote distribution***Comments****17Master** **Highly Voted** 2 years, 7 months ago**Selected Answer: B**

Tricky question!!! forget one-way or two-way. In this scenario, AWS applications (Amazon Chime, Amazon Connect, Amazon QuickSight, AWS Single Sign-On, Amazon WorkDocs, Amazon WorkMail, Amazon WorkSpaces, AWS Client VPN, AWS Management Console, and AWS Transfer Family) need to be able to look up objects from the on-premises domain in order for them to function. This tells you that authentication needs to flow both ways. This scenario requires a two-way trust between the on-premises and AWS Managed Microsoft AD domains.

It is a requirement of the application

Scenario 2: <https://aws.amazon.com/es/blogs/security/everything-you-wanted-to-know-about-trusts-with-aws-managed-microsoft-ad/>

upvoted 78 times

mohamedsambo 1 year, 5 months ago

AWS IAM Identity Center requires a two-way trust so that it has permissions to read user and group information from your domain to synchronize user and group metadata. IAM Identity Center uses this metadata when assigning access to permission sets or applications. User and group metadata is also used by applications for collaboration, like when you share a dashboard with another user or group. The trust from AWS Directory Service for Microsoft Active Directory to your domain permits IAM Identity Center to trust your domain for authentication. The trust in the opposite direction grants AWS permissions to read user and group metadata.

upvoted 7 times

aatikah 6 months ago

Application Context in This Scenario:
AWS Applications Mentioned in the Scenario:

The question focuses on AWS Single Sign-On (AWS SSO) and the use of SSO across AWS accounts.
AWS SSO with AWS Organizations does not require AWS applications to look up or access on-premises AD objects beyond authentication.
Blog Reference:

The blog mentions the need for a two-way trust if AWS services like Amazon WorkSpaces or Amazon WorkMail require access to user attributes stored in the on-premises AD.
This scenario does not mention such AWS applications, so a one-way trust is sufficient for the SSO use case.

upvoted 2 times

pbpally 2 years, 1 month ago

The problem with this is that nowhere in the question is it saying that the application needs to be able to flow back so two-way is not needed.

upvoted 4 times

pbpally 2 years, 1 month ago

What I did find though was documentation that explicitly states that IAM Identity Center (successor to AWS SSO) requires a two-way trust:

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/ms_ad_setup_trust.html

upvoted 12 times

KADSM Highly Voted 2 years, 7 months ago

Answer B as we have AWS SSO which requires two way trust. As per documentation - A two-way trust is required for AWS Enterprise Apps such as Amazon Chime, Amazon Connect, Amazon QuickSight, AWS IAM Identity Center (successor to AWS Single Sign-On), Amazon WorkDocs, Amazon WorkMail, Amazon WorkSpaces, and the AWS Management Console. AWS Managed Microsoft AD must be able to query the users and groups in your self-managed AD.

Amazon EC2, Amazon RDS, and Amazon FSx will work with either a one-way or two-way trust.

upvoted 14 times

pbpally 2 years, 1 month ago

I found the documentation that explicitly states that IAM Identity Center (successor to AWS SSO) requires a two-way trust:

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/ms_ad_setup_trust.html

upvoted 6 times

asad_devops_911 Most Recent 1 month, 1 week ago

Selected Answer: A

One-way trust allows users from the on-prem AD to authenticate into AWS resources via IAM Identity Center, without granting AWS access to manage the on-prem AD.

upvoted 1 times

Juju66 1 month, 3 weeks ago

Selected Answer: A

I think one-way trust is sufficient for AWS SSO to authenticate users from the on-prem AD without requiring changes to the on-prem directory. two-way trust is unnecessary for SSO and introduces additional complexity.

Both two-way and one-way are valid for forest trust.

[https://learn.microsoft.com/en-us/previous-versions/windows/it-pro/windows-server-2008-r2-and-2008/cc771397\(v=ws.10\)](https://learn.microsoft.com/en-us/previous-versions/windows/it-pro/windows-server-2008-r2-and-2008/cc771397(v=ws.10))

upvoted 1 times

omunoz 3 months ago

Selected Answer: A

A one-way trust (rather than a two-way trust) is the recommended approach for security reasons:
The on-premises AD remains the authoritative identity provider.
AWS trusts the on-prem AD, but AD does not trust AWS.
This setup ensures that AWS does not introduce potential security risks to the on-premises

A two-way trust is unnecessary for SSO and can introduce security risks, as it allows AWS to authenticate users in the on-prem AD and vice versa.

AWS best practices recommend using a one-way trust for such integrations.

upvoted 3 times

javaguy1906 3 months, 3 weeks ago

Selected Answer: A

A one-way trust is sufficient because AWS should trust the on-premises AD for authentication, but the on-premises AD does not need to trust AWS

upvoted 1 times

AwsAbhiKumar 4 months, 2 weeks ago

Selected Answer: A

AWS Single Sign-On can integrate with on-premises Microsoft Active Directory using AWS Directory Service for Microsoft Active Directory.

A one-way trust relationship is sufficient for AWS SSO to authenticate users from the on-premises Active Directory without requiring changes to existing trust relationships in the on-premises environment.

upvoted 1 times

AfricanCloudGuru 4 months, 2 weeks ago

Selected Answer: B

Because we need to create two ways forest traffic

upvoted 1 times

Rcosmos 5 months ago

Selected Answer: B

Resposta correta: B!

Explicação revisada:

Quando a empresa precisa integrar seu Microsoft Active Directory autogerenciado com o AWS SSO para login único em múltiplas contas na AWS, uma relação de confiança de floresta bidirecional é necessária para permitir que usuários e grupos locais sejam autenticados e tenham acesso completo aos recursos da AWS. Essa abordagem garante interoperabilidade total entre o AD local e o AWS SSO.

upvoted 1 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: A

AWS Single Sign-On (AWS SSO) allows centralized management of access to multiple AWS accounts, and it integrates with on-premises Active Directory via AWS Directory Service for Microsoft Active Directory.

A one-way trust enables AWS Directory Service to authenticate users from the on-premises AD without giving AWS access to manage the on-premises AD. This is secure and follows best practices.

This is the correct solution because it integrates the on-premises AD with AWS SSO and maintains the centralized management of users and groups in the on-premises AD.

upvoted 3 times

John54321 5 months, 2 weeks ago

Selected Answer: B

AWS SSO (previously called Identity center) requires two way trust.

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/ms_ad_setup_trust.html

upvoted 1 times

hilker1983 5 months, 3 weeks ago

Selected Answer: A

AWS Single Sign-On (AWS SSO): Provides a centralized way to manage SSO access to multiple AWS accounts and business applications. It simplifies user access management across AWS Organizations.

One-Way Trust: Creating a one-way trust with AWS Directory Service allows the on-premises Microsoft Active Directory to authenticate users without allowing AWS to initiate authentication requests to the on-premises directory. This maintains security and control.

Minimal Operational Overhead: This approach leverages AWS managed services, reducing the complexity and overhead of managing SSO and directory services.

upvoted 1 times

aatikah 6 months ago

Selected Answer: A

Application Context in This Scenario:
AWS Applications Mentioned in the Scenario:

The question focuses on AWS Single Sign-On (AWS SSO) and the use of SSO across AWS accounts. AWS SSO with AWS Organizations does not require AWS applications to look up or access on-premises AD objects beyond authentication.

This scenario does not mention such AWS applications, so a one-way trust is sufficient for the SSO use case.

upvoted 1 times

Carlini2020 7 months ago

Selected Answer: A

No need for both ways communication.

upvoted 1 times

trinh_le 7 months ago

Selected Answer: B

A specific aws applications require two-way trust. To migrate to the aws, we need all apps are available. So I pick B

upvoted 1 times

mzeynalli 7 months, 3 weeks ago

Selected Answer: A

A. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console. Create a one-way forest trust or a one-way domain trust to connect the company's self-managed Microsoft Active Directory with AWS SSO by using AWS Directory Service for Microsoft Active Directory.

Two-way trust (as mentioned in option B) allows mutual trust between both environments, which is not necessary if the company only wants AWS to authenticate against the on-premises AD.

A two-way trust is not necessary if the company only needs authentication from the on-premises AD. It may introduce unnecessary complexity and potential security risks since it requires trust in both directions.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Current version of this question (with Identity Center) doesn't even contain option A.

<https://www.examttopics.com/discussions/amazon/view/136806-exam-aws-certified-solutions-architect-associate-saa-c03/>

upvoted 5 times

techghost 5 months, 1 week ago

Looks like this tell us that B is the answer

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #29

Topic 1

A company provides a Voice over Internet Protocol (VoIP) service that uses UDP connections. The service consists of Amazon EC2 instances that run in an Auto Scaling group. The company has deployments across multiple AWS Regions. The company needs to route users to the Region with the lowest latency. The company also needs automated failover between Regions.

Which solution will meet these requirements?

- A. Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Use the NLB as an AWS Global Accelerator endpoint in each Region. **Most Voted**
- B. Deploy an Application Load Balancer (ALB) and an associated target group. Associate the target group with the Auto Scaling group. Use the ALB as an AWS Global Accelerator endpoint in each Region.
- C. Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Create an Amazon Route 53 latency record that points to aliases for each NLB. Create an Amazon CloudFront distribution that uses the latency record as an origin.
- D. Deploy an Application Load Balancer (ALB) and an associated target group. Associate the target group with the Auto Scaling group. Create an Amazon Route 53 weighted record that points to aliases for each ALB. Deploy an Amazon CloudFront distribution that uses the weighted record as an origin.

Correct Answer: A*Community vote distribution***Comments****Six_Fingered_Jose** **Highly Voted** 2 years, 7 months ago**Selected Answer: A**

agree with A,
Global Accelerator has automatic failover and is perfect for this scenario with VoIP
<https://aws.amazon.com/global-accelerator/faqs/>
upvoted 61 times

awashenko 1 year, 8 months ago

I also agree A after reading this link.

upvoted 4 times

bnagaraja9099 1 year, 7 months ago

A - Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover. Both services integrate with AWS Shield for DDoS protection.

upvoted 10 times

TilTil 1 year, 3 months ago

This is the best case for A to be an answer. Cloudfront is great but for HTTP use cases.

upvoted 2 times

BoboChow 2 years, 7 months ago

Thank you for your link, it make me consolidate A.

upvoted 7 times

bullrem 2 years, 4 months ago

This option does not meet the requirements because AWS Global Accelerator is only used to route traffic to the optimal AWS Region, it does not provide automatic failover between regions.

upvoted 2 times

sachin 2 years, 3 months ago

Instant regional failover: AWS Global Accelerator automatically checks the health of your applications and routes user traffic only to healthy application endpoints. If the health status changes or you make configuration updates, AWS Global Accelerator reacts instantaneously to route your users to the next available endpoint.

upvoted 10 times

ElaineRan 1 year, 10 months ago

Thank you, the link also helps me to know the differences between Global Acc and CloudFront.

upvoted 3 times

mouhannadhaj Highly Voted 2 years, 7 months ago

Selected Answer: A

CloudFront uses Edge Locations to cache content while Global Accelerator uses Edge Locations to find an optimal pathway to the nearest regional endpoint. CloudFront is designed to handle HTTP protocol meanwhile Global Accelerator is best used for both HTTP and non-HTTP protocols such as TCP and UDP. so i think A is a better answer

upvoted 40 times

LEE0scar Most Recent 4 weeks ago

Selected Answer: A

UDP Support:

UDP traffic (used by VoIP) is not supported by Application Load Balancer (ALB).

Network Load Balancer (NLB) does support UDP, making NLB the correct choice.

Lowest latency routing:

AWS Global Accelerator is designed to route traffic to the AWS Region with the lowest latency.

It uses the AWS global network for optimal performance, which is better than Route 53 latency-based routing.

upvoted 3 times

Juju66 1 month, 3 weeks ago

Selected Answer: A

UDP - User Datagram Protocol (UDP) is a transport layer protocol used for sending messages (datagrams) between applications over a network. UDP is part of the Internet Protocol (IP) suite and is defined by Internet Engineering Task Force (IETF) standards in RFC 768.

upvoted 1 times

lubomir.tomecek 4 months, 1 week ago

Selected Answer: A

C is not correct, because Network Loadbalancer is already integrated with Route 53.

upvoted 1 times

adamatic 4 months, 1 week ago

Selected Answer: A

i full understand this is a use case for global accelerator but can anyone tell me why you need a NLB over an ALB here?

upvoted 1 times

adamatic 4 months, 1 week ago

oh i see, ALB is for HTTP/HTTPS

upvoted 2 times

AfricanCloudGuru 4 months, 2 weeks ago

Selected Answer: A

The Nature of the application VoIP needs NLB and Global Accelerator + an auto scalling

upvoted 1 times

ylanzzz 5 months, 1 week ago

Selected Answer: A

cloudfront doesn't provide automatic failover

upvoted 1 times

satyaamm 5 months, 2 weeks ago

Selected Answer: A

NLB is suitable for UDP and global accelerator is also suitable for UDP.

upvoted 1 times

AbhishekCloudEngg 5 months, 2 weeks ago

Selected Answer: A

The correct answer is A: Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Use the NLB as an AWS Global Accelerator endpoint in each Region.

Key reasons:

NLB supports UDP protocol (required for VoIP), while ALB only handles HTTP/HTTPS

Global Accelerator provides:

Automatic routing to lowest latency region

Built-in cross-region failover

Optimized network performance via AWS global network

upvoted 1 times

minhhieu_tech_1 6 months ago

Selected Answer: A

Cloudfront is not support UDP Protocol. So maybe C answer isn't correct.

upvoted 1 times

EzKkk 6 months, 3 weeks ago

Selected Answer: C

I think this is not a very well written one. What we need:

1, Network Service that handles latency based policy and cross region failover -> Route53

2, A component that route traffic to ASG that are compatible with UDP -> Network Load Balancer.

upvoted 1 times

Bobby3 10 months, 1 week ago

Selected Answer: C

Aws GA is most suitable for https.

upvoted 1 times

hichtr2 11 months ago

prashant 11 months ago

A : UDP so NLB and Global Accelerator reduces the number of hops by providing packets to travel over congestion free AWS global network.

Global Accelerator supported end point : ALB,NLB,EC2 & Elastic IP address

upvoted 2 times

jatric 11 months, 2 weeks ago

Selected Answer: A

you can configure a Network Load Balancer (NLB) in each AWS Region to address your on-premises endpoints. Then you can register the NLBs as endpoints in your AWS Global Accelerator configuration.

<https://aws.amazon.com/global-accelerator/faqs/>

upvoted 2 times

ManikRoy 1 year, 2 months ago

Selected Answer: A

UDP Connection :- So NLB

Routing to region having lowest latency and also with Automated failover, Also non-HTTP use cases, such as gaming (UDP), or

Voice over IP - Global Accelerator

upvoted 2 times

TruthWS 1 year, 2 months ago

A is correct because Accelerator endpoint usefull more than route53

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #30

Topic 1


A development team runs monthly resource-intensive tests on its general purpose Amazon RDS for MySQL DB instance with Performance Insights enabled. The testing lasts for 48 hours once a month and is the only process that uses the database. The team wants to reduce the cost of running the tests without reducing the compute and memory attributes of the DB instance. Which solution meets these requirements MOST cost-effectively?

- A. Stop the DB instance when tests are completed. Restart the DB instance when required.
- B. Use an Auto Scaling policy with the DB instance to automatically scale when tests are completed.
- C. Create a snapshot when tests are completed. Terminate the DB instance and restore the snapshot when required.
- D. Modify the DB instance to a low-capacity instance when tests are completed. Modify the DB instance again when required.

Most Voted

Correct Answer: C

Community vote distribution



Comments

hanhdroid **Highly Voted** 2 years, 8 months ago

Selected Answer: C

Answer C, you still pay for storage when an RDS database is stopped
upvoted 42 times

KVK16 **Highly Voted** 2 years, 8 months ago

Selected Answer: C

C - Create a manual Snapshot of DB and shift to S3- Standard and Restore from Manual Snapshot when required.

Not A - By stopping the DB although you are not paying for DB hours you are still paying for Provisioned IOPs, the storage for Stopped DB is more than Snapshot of underlying EBS vol. and Automated Back ups.

Not D - Is possible but not MOST cost effective, no need to run the RDS when not needed.

upvoted 15 times

Juju66 **Most Recent** 1 month, 3 weeks ago

Selected Answer: C

Stopping an RDS instance does not reduce costs significantly because you still pay for the storage and some management overhead. Additionally, stopping and starting an RDS instance can take time and may not be as efficient as terminating and restoring from a snapshot.

upvoted 2 times

donquixote47 2 months ago

Selected Answer: C

you still pay for storage when an RDS database is stopped

upvoted 1 times

Kazmin 2 months, 2 weeks ago

Selected Answer: B

B is right i think

upvoted 1 times

John54321 5 months, 2 weeks ago

Selected Answer: C

An RDS instance can be stopped for up to seven days so it cannot be A.

"The instance stops running, up to a maximum of 7 consecutive days."

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_StopInstance.html

upvoted 3 times

f51a8bd 5 months, 3 weeks ago

Selected Answer: A

La opción A es la más rentable porque detiene la instancia de base de datos cuando no se está utilizando, ahorrando los costos de cómputo, y luego se reinicia cuando sea necesario, sin comprometer el rendimiento durante las pruebas.

upvoted 1 times

mzeynalli 7 months, 3 weeks ago

Selected Answer: C

The correct answer is:

C. Create a snapshot when tests are completed. Terminate the DB instance and restore the snapshot when required.

Explanation:

This solution is the most cost-effective because it allows the development team to:

Save Costs by Terminating the Instance:

When the DB instance is not in use (28 out of 30 days), terminating it eliminates the costs associated with running a live instance.

You only pay for storage costs associated with the snapshot, which is significantly lower than the cost of keeping a running instance.

Restoring a Snapshot:

Restoring from a snapshot is straightforward and brings the database back with the same data and configuration as before.

It's a good practice when dealing with periodic, resource-intensive tasks like the monthly tests.

upvoted 2 times

PaulGa 10 months ago

Selected Answer: C

Ans C – take Snapshots and restore them, because otherwise you're still paying for RDS storage.

upvoted 2 times

OBIOHAnze 1 year ago

Selected Answer: C

C. This option allows you to save on costs by only paying for storage of the snapshot when the DB instance is terminated. When needed again, you can restore the DB instance from the snapshot, which is a cost-effective way to handle infrequent but resource-intensive tasks.

upvoted 6 times

ManikRoy 1 year, 1 month ago

Selected Answer: C

Most cost effective is to create a snapshot and get rid of the DB instance after testing.

Note that A is not correct option as While your database instance is stopped, you are charged for provisioned storage, manual snapshots and automated backup storage within your specified retention window, but not for database instance hours.

upvoted 2 times

AWSCLOUDLMD 1 year, 1 month ago

Selected Answer: A

A Since the tests only run once a month for 48 hours, this approach minimizes costs while still retaining the same compute and memory attributes when the instance is restarted.

when Snapshot resorts i new config

upvoted 4 times

vi24 1 year, 3 months ago

My question is: isn't this DB collecting new data during the testing period (48 hrs.) ? after the snapshot is taken ? stop and restore db from the snapshot is the most cost effective but I think some data might be lost in between, so wouldn't be feasible !

upvoted 2 times

VanDacker 1 year, 3 months ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_StopInstance.html

upvoted 5 times

FlyingHawk 5 months, 1 week ago

Based on the AWS RDS documentation, an RDS instance can only be stopped for a maximum of 7 days. After 7 days, the instance is automatically started by AWS to ensure data integrity and maintenance.

upvoted 1 times

Isomas 1 year, 3 months ago

Answer C because,

You can stop a DB instance for up to seven days. If you don't manually start your DB instance after seven days, your DB instance is automatically started so that it doesn't fall behind any required maintenance updates.

So, the "auto starting" behaviour is expected.

If you rarely use the database, BEST option is to Snapshot and Delete the database. Then, when you need it again, you could launch a new database from the Snapshot.

Amazon RDS is not intended to be stopped for long periods.

upvoted 4 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

DB is used one a month for 48 hours only so there is no point in keeping it up for rest of the month.

B: more cost

D: not allowed to reduce computing power

A: It will work but C is much cheaper as instance is not only stopped but terminated.

upvoted 2 times

A_jaa 1 year, 5 months ago

Selected Answer: C

Answer-C

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #31

Topic 1

A company that hosts its web application on AWS wants to ensure all Amazon EC2 instances, Amazon RDS DB instances, and Amazon Redshift clusters are configured with tags. The company wants to minimize the effort of configuring and operating this check.

What should a solutions architect do to accomplish this?

- A. Use AWS Config rules to define and detect resources that are not properly tagged. **Most Voted**
- B. Use Cost Explorer to display resources that are not properly tagged. Tag those resources manually.
- C. Write API calls to check all resources for proper tag allocation. Periodically run the code on an EC2 instance.
- D. Write API calls to check all resources for proper tag allocation. Schedule an AWS Lambda function through Amazon CloudWatch to periodically run the code.

Correct Answer: A

Community vote distribution

A (100%)

Comments

kurinei021 **Highly Voted** 2 years, 5 months ago

Answer from ChatGPT:

Yes, you can use AWS Config to create tags for your resources. AWS Config is a service that enables you to assess, audit, and evaluate the configurations of your AWS resources. You can use AWS Config to create rules that automatically tag resources when they are created or when their configurations change.

To create tags for your resources using AWS Config, you will need to create an AWS Config rule that specifies the tag key and value you want to use and the resources you want to apply the tag to. You can then enable the rule and AWS Config will automatically apply the tag to the specified resources when they are created or when their configurations change.

upvoted 22 times

aaroncelestin 1 year, 10 months ago

This the first answer that I've seen ChatGPT get correct here on ExamTopics. You should all know that using ChatGPT for this is bound to give bad answers. It only parrots what it has seen written/copied/pasted by someone/something somewhere, picked up with absolutely zero context. ChatGPT doesn't "know" anything about AWS services. So, beware the "answers" it gives.

upvoted 14 times

kidomaruto 1 year, 7 months ago

I tried it with Bing AI, and the answer was almost always the right one.
It depends a lot on the prompt quality
upvoted 3 times

cookieMr Highly Voted 1 year, 12 months ago

Selected Answer: A

AWS Config provides a set of pre-built or customizable rules that can be used to check the configuration and compliance of AWS resources. By creating a custom rule or using the built-in rule for tagging, you can define the required tags for EC2, RDS DB and Redshift clusters. AWS Config continuously monitors the resources and generates configuration change events or evaluation results.

By leveraging AWS Config, the solution can automatically detect any resources that do not comply with the defined tagging requirements. This approach eliminates the need for manual checks or periodic code execution, reducing operational overhead. Additionally, AWS Config provides the ability to automatically remediate non-compliant resources by triggering Lambda or sending notifications, further streamlining the configuration management process.

Option B (using Cost Explorer) primarily focuses on cost analysis and does not provide direct enforcement of proper tagging. Option C and D (writing API calls and running them manually or through scheduled Lambda) require more manual effort and maintenance compared to using AWS Config rules.

upvoted 10 times

AshishDhole Most Recent 5 months ago

Selected Answer: A

By Using AWS Config we can tag the resources, and then filter the resources which are tagged or not tagged.

upvoted 1 times

satyaamm 5 months, 2 weeks ago

Selected Answer: A

AWS Config helps identify any configuration issues.

upvoted 1 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: A

AWS Config is designed to monitor and assess the compliance of AWS resource configurations.

You can use AWS Config managed rules, such as required-tags, to automatically evaluate whether resources are properly tagged.

AWS Config provides a centralized dashboard to view non-compliant resources and requires minimal setup and maintenance. This is the correct answer because it provides an automated, scalable, and low-effort solution to detect resources that are not properly tagged.

upvoted 1 times

mzeynalli 7 months, 3 weeks ago

Selected Answer: A

Answer: A. Use AWS Config rules to define and detect resources that are not properly tagged.

Explanation:

Using AWS Config is the best solution because it provides a managed service to monitor and evaluate the configurations of your AWS resources continuously. Here's why AWS Config is the preferred choice:

Automated Tag Compliance:

Ease of Setup and Operation:

Comprehensive Coverage:

AWS Config can monitor a wide range of AWS resources, including Amazon EC2 instances, RDS databases, Redshift clusters, and more, ensuring that all your resources are properly tagged.

upvoted 3 times

PaulGa 9 months, 2 weeks ago

Selected Answer: A

Not sure we need ChatGPT here (other than to prove what it can/can't do): answers B, C, D all require manual - and periodic - intervention. Has to be Ans A.

upvoted 1 times

DavidNgTan 10 months, 2 weeks ago

Selected Answer: A

You can use AWS Config to create a rule that evaluates whether your resources have the required tags applied.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

<https://docs.aws.amazon.com/whitepapers/latest/tagging-best-practices/implementing-and-enforcing-tagging.html>

AWS Config (required_tag)

AWS Config is a service that allows you to assess, audit, and evaluate the configurations of your AWS resources (see Resource types supported by AWS Config). In the case of tagging, we can use it to identify resources that are lacking tags with specific keys, using the required_tags rule (refer to Resource types supported by required_tags). From the earlier example, we might test for the existence of the key on all Amazon EC2 instances.

upvoted 3 times

A_jaa 1 year, 5 months ago

Selected Answer: A

Answer-A

upvoted 1 times

Ruffyt 1 year, 7 months ago

Has typos in the question, correct is "A company that hosts its web application on AWS wants to ensure all Amazon EC2 instance, Amazon RDS DB instances, and Amazon Redshift clusters are configured with tags." Keyword "are configured with tags", choose (A) "AWS Config rules".

upvoted 2 times

awashenko 1 year, 8 months ago

Selected Answer: A

I originally thought D, but after reading through the discussion I agree that option A would require less effort. D would get the job done but would require more effort so I think A is correct.

upvoted 2 times

KawtarZ 1 year, 9 months ago

Selected Answer: A

A without a doubt

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: A

AWS Config continually assesses, audits, and evaluates the configurations and relationships of your resources on AWS, on premises, and on other clouds.

upvoted 3 times

james2033 1 year, 10 months ago

Selected Answer: A

Has typos in the question, correct is "A company that hosts its web application on AWS wants to ensure all Amazon EC2 instance, Amazon RDS DB instances, and Amazon Redshift clusters are configured with tags." Keyword "are configured with tags", choose (A) "AWS Config rules".

upvoted 2 times

miki111 1 year, 11 months ago

Option A is the right answer for this.

upvoted 1 times

lelouchjedai 1 year, 12 months ago

Selected Answer: A

The answer is A
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #32

Topic 1

A development team needs to host a website that will be accessed by other teams. The website contents consist of HTML, CSS, client-side JavaScript, and images.

Which method is the MOST cost-effective for hosting the website?

- A. Containerize the website and host it in AWS Fargate.
- B. Create an Amazon S3 bucket and host the website there. **Most Voted**
- C. Deploy a web server on an Amazon EC2 instance to host the website.
- D. Configure an Application Load Balancer with an AWS Lambda target that uses the Express.js framework.

Correct Answer: B

Community vote distribution

B (100%)

Comments

masetromain **Highly Voted** 2 years, 8 months ago

Selected Answer: B

Good answer is B: client-side JavaScript. the website is static, so it must be S3.

upvoted 35 times

BoboChow **Highly Voted** 2 years, 8 months ago

Selected Answer: B

HTML, CSS, client-side JavaScript, and images are all static resources.

upvoted 10 times

dattateja8 **Most Recent** 2 months, 2 weeks ago

Selected Answer: B

Ans B: Don't get confused by looking at JavaScript which can be dynamic. It is Client-side JavaScript. Means it get's executed at the Client side in a web browser. So it is all static and S3 is the cheapest option to host a static website.

upvoted 1 times

satyaamm 5 months, 2 weeks ago

Selected Answer: B

S3 is the most suitable option for hosting static websites.

upvoted 1 times

PaulGa 9 months, 2 weeks ago

Selected Answer: B

I had initially thought Ans A... but its Ans C -- "cookieMr" makes it clear:

Ans A "Containerising the website and hosting with AWS Fargate involves additional complexity and costs associated with managing the container environment and scaling resources."

So it has to be...

Ans B: "...Amazon S3 to host the website, take advantage of its durability, scalability, and low-cost pricing model. Only pay for the storage and data transfer associated with your website, without the need for managing and maintaining web servers or containers. This reduces the operational overhead and infrastructure costs."

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

Cheapest Static site hosting = S3

upvoted 3 times

A_jaa 1 year, 5 months ago

Selected Answer: B

Answer-B

upvoted 2 times

Ruffyit 1 year, 7 months ago

HTML, CSS, client-side JavaScript, and images are all static resources.

upvoted 1 times

AWSStudyBuddy 1 year, 7 months ago

The MOST cost-effective method for hosting a website is to:

Create an Amazon S3 bucket and host the website there.

Amazon S3 is a highly scalable and cost-effective object storage service. It is a good option for hosting static websites, such as the website in this scenario.

To host a static website on Amazon S3, you would first need to create an S3 bucket. Then, you would need to upload the website files to the bucket. Once the files are uploaded, you can configure the bucket to serve as a website.

upvoted 2 times

hungpm 1 year, 9 months ago

Selected Answer: B

Static website should work fine with S3

upvoted 1 times

KawtarZ 1 year, 9 months ago

Selected Answer: B

the website is static because the backend runs on client side.

upvoted 2 times

evanhongo 1 year, 10 months ago

Selected Answer: B

all static resources.

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: B

static website, cost-effective = S3 web hosting

upvoted 3 times

james2033 1 year, 10 months ago

Selected Answer: B

Just all static content HTML, CSS, client-side JavaScript, images. Amazon S3 is good enough.

upvoted 1 times

miki111 1 year, 11 months ago

Option B is the right answer for this.

upvoted 1 times

Kaab_B 1 year, 11 months ago

Selected Answer: B

S3 is amongst the cheapest services offered by AWS.

upvoted 1 times

karloscetina007 1 year, 11 months ago

Selected Answer: B

B is the correct answer.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #33

Topic 1


A company runs an online marketplace web application on AWS. The application serves hundreds of thousands of users during peak hours. The company needs a scalable, near-real-time solution to share the details of millions of financial transactions with several other internal applications. Transactions also need to be processed to remove sensitive data before being stored in a document database for low-latency retrieval.

What should a solutions architect recommend to meet these requirements?

- A. Store the transactions data into Amazon DynamoDB. Set up a rule in DynamoDB to remove sensitive data from every transaction upon write. Use DynamoDB Streams to share the transactions data with other applications.
- B. Stream the transactions data into Amazon Kinesis Data Firehose to store data in Amazon DynamoDB and Amazon S3. Use AWS Lambda integration with Kinesis Data Firehose to remove sensitive data. Other applications can consume the data stored in Amazon S3.
- C. Stream the transactions data into Amazon Kinesis Data Streams. Use AWS Lambda integration to remove sensitive data from every transaction and then store the transactions data in Amazon DynamoDB. Other applications can consume the transactions data off the Kinesis data stream. **Most Voted**
- D. Store the batched transactions data in Amazon S3 as files. Use AWS Lambda to process every file and remove sensitive data before updating the files in Amazon S3. The Lambda function then stores the data in Amazon DynamoDB. Other applications can consume transaction files stored in Amazon S3.

Correct Answer: C

Community vote distribution



Comments

ArielSchivo **Highly Voted** 2 years, 7 months ago

Selected Answer: C

I would go for C. The tricky phrase is "near-real-time solution", pointing to Firehose, but it can't send data to DynamoDB, so it leaves us with C as best option.

Kinesis Data Firehose currently supports Amazon S3, Amazon Redshift, Amazon OpenSearch Service, Splunk, Datadog, NewRelic, Dynatrace, Sumologic, LogicMonitor, MongoDB, and HTTP End Point as destinations.

https://aws.amazon.com/kinesis/data-firehose/#_h1=Kinesis%20Data%20Firehose%20currently%20supports%20HTTP%20End%20Point%20destinations

firehose/taqs/#:~:text=Kinesis%20Data%20Firehose%20currently%20supports,H I F%20End%20Point%20as%20destinations.
upvoted 101 times

Lonojack 2 years, 4 months ago

This was a really tough one. But you have the best explanation on here with reference point. Thanks. I'm going with answer C!
upvoted 4 times

SaraSundaram 2 years, 2 months ago

There are many questions having Firehose and Stream. Need to know them in detail to answer. Thanks for the explanation
upvoted 4 times

diabloexodia 1 year, 11 months ago

Stream is used if you want real time results , but with firehose , you generally use the data at a later point of time by storing it somewhere. Hence if you see "REAL TIME" the answer is most probably Kinesis Data Streams.
upvoted 20 times

lizzard812 2 years, 4 months ago

Sorry but I still can't see how Kinesis Data Stream is 'scalable', since you have to provision the quantity of shards in advance?
upvoted 1 times

habibi03336 2 years, 3 months ago

"easily stream data at any scale"
This is a description of Kinesis Data Stream. I think you can configure its quantity but still not provision and manage scalability by yourself.
upvoted 1 times

JesseeS **Highly Voted** 2 years, 7 months ago

The answer is C, because Firehose does not support DynamoDB and another key word is "data" Kinesis Data Streams is the correct choice. Pay attention to key words. AWS likes to trick you up to make sure you know the services.
upvoted 33 times

kyd0nix **Most Recent** 4 months, 1 week ago

Selected Answer: B

IMO "near-real-time" is key for Firehose, BUT since of all the discussions B vs C (Firehose can't have DynamoDB as destination, I think the question is misswritten and has to be reviewed to avoid the confusion)
upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: C

C - Kinesis Data Streams allows for low-latency processing, which is crucial for near-real-time requirements.
upvoted 1 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: C

Scalable processing: The system must scale to handle hundreds of thousands of users and millions of transactions during peak hours.
Near-real-time sharing of transactions: Transactions should be shared with internal applications in near-real time.
Sensitive data removal: Sensitive information must be processed and removed before storage.
Low-latency retrieval: The processed data must be stored in a document database (Amazon DynamoDB) for quick access.
upvoted 1 times

aefuen1 5 months, 3 weeks ago

Selected Answer: B

It's B. You can write to the DynamoDB table from the lambda preprocessing function. Also option C can't be correct because if "Other applications can consume the transactions data off the Kinesis data stream", this means they will consume data with sensitive values, which is a constraint for the solution.
upvoted 2 times

kamel1 3 months, 4 weeks ago

kehen 9 months, 4 weeks ago

The question says it needs to have sensitive information removed before DB storage. Other apps could consume transaction data but not necessarily store, or consume only non-sensitive data.

upvoted 1 times

engnrshon 6 months, 3 weeks ago

C :

upvoted 1 times

Mauro0001 9 months, 1 week ago

Selected Answer: C

One of the tricky phrases is 'near-real-time solutions' because it points to the fact that every time a write is made to a database, it incurs a delay, and then retrieving it with an API call adds another latency. With Kinesis Data Streams, that process is optimized because the intermediary that gives you the ability to write to DynamoDB also provides that data to other services due to the retention period of Kinesis Data Streams.

upvoted 2 times

PaulGa 9 months, 2 weeks ago

Selected Answer: C

Ans C.

High level difference between the Kinesis and DynamoDB:

Kinesis Streams allows production/ consumption of large volumes of data (web data, logs, etc); DynamoDB Streams is a feature local to DynamoDB to track the granular changes to DynamoDB table items.

(Note also: data latency for Firehose is 60 seconds or higher; Streams is for custom processing and has sub-second processing latency).

upvoted 2 times

Lin878 11 months, 1 week ago

Selected Answer: C

Q: What is a destination in Firehose?

A destination is the data store where your data will be delivered. Firehose currently supports Amazon S3, Amazon Redshift, Amazon OpenSearch Service, Splunk, Datadog, NewRelic, Dynatrace, Sumo Logic, LogicMonitor, MongoDB, and HTTP End Point as destinations.

<https://aws.amazon.com/firehose/faqs/>

upvoted 2 times

the_mellie 1 year ago

Selected Answer: C

with multiple consumers and on the fly modification, it seems like the most logical choice

upvoted 2 times

vi24 1 year, 3 months ago

I chose B. The "near real time" is very specific to Kinesis firehose which is a better option anyway. The rest of the answer makes sense too. C is wrong : "sensitive data removed by Lambda & then store transaction data in DynamoDB" , while it continues to say other applications are accessing the transaction data from kinesis Data stream !!

upvoted 3 times

Pics00094 1 year, 3 months ago

Selected Answer: C

need to know..

1) Lambda Integration

2) Difference between Real time(Kinesis Data Stream) vs Near Real time(Kinesis Fire House)

3) Firehouse can't target DynamoDB

upvoted 5 times

JulianWaksmann 1 year, 4 months ago

i think c are bad too, because it isn't near real time.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

A: DynamoDB streams are logs, not fit for real-time sharing.

B: S3 is not document database, it's BLOB

D: S3 and files are not database

C: Kinesis + Lambda + DynamoDB is high performance, low latency scalable solution.

upvoted 3 times

A_jaa 1 year, 5 months ago

Selected Answer: C

Answer-C

upvoted 1 times

bujuman 1 year, 5 months ago

Selected Answer: C

Data Stream can handle near-real-time and is able to store to DynamoDB

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #34

Topic 1

A company hosts its multi-tier applications on AWS. For compliance, governance, auditing, and security, the company must track configuration changes on its AWS resources and record a history of API calls made to these resources. What should a solutions architect do to meet these requirements?

- A. Use AWS CloudTrail to track configuration changes and AWS Config to record API calls.
- B. Use AWS Config to track configuration changes and AWS CloudTrail to record API calls. **Most Voted**
- C. Use AWS Config to track configuration changes and Amazon CloudWatch to record API calls.
- D. Use AWS CloudTrail to track configuration changes and Amazon CloudWatch to record API calls.

Correct Answer: B

Community vote distribution

B (99%)

A

Comments

airraid2010 **Highly Voted** 2 years, 7 months ago

Selected Answer: B

CloudTrail - Track user activity and API call history.
Config - Assess, audits, and evaluates the configuration and relationships of tag resources.

Therefore, the answer is B

upvoted 40 times

satyaamm **Most Recent** 5 months, 2 weeks ago

Selected Answer: B

AWS Config for configuration changes and AWS CloudTrail for API activities.
upvoted 2 times

SBS2 6 months, 3 weeks ago

Selected Answer: B

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

Config = Governance, auditing of AWS resource
CloudTrail = API call tracking
B is correct

upvoted 4 times

A_jaa 1 year, 5 months ago

Selected Answer: B

Answer-B

upvoted 1 times

Ruffyt 1 year, 7 months ago

Correct Answer- Option B. Here's why

AWS Config for Configuration Changes: AWS Config is a service that tracks changes to resource configurations over time. It provides a history of configuration changes to your AWS resources and helps with compliance and auditing by allowing you to assess how resource configurations have changed over time.

AWS CloudTrail for API Calls: AWS CloudTrail is designed specifically for recording API calls made to AWS resources. It captures detailed information about who made each API call, the actions taken, and the resources affected. This is essential for auditing and security purposes.

upvoted 3 times

AWSStudyBuddy 1 year, 7 months ago

Correct Answer- Option B. Here's why

AWS Config for Configuration Changes: AWS Config is a service that tracks changes to resource configurations over time. It provides a history of configuration changes to your AWS resources and helps with compliance and auditing by allowing you to assess how resource configurations have changed over time.

AWS CloudTrail for API Calls: AWS CloudTrail is designed specifically for recording API calls made to AWS resources. It captures detailed information about who made each API call, the actions taken, and the resources affected. This is essential for auditing and security purposes.

While Amazon CloudWatch can be used to monitor and gather metrics, it is not designed for recording API calls or tracking configuration changes. AWS Config and AWS CloudTrail are purpose-built for these specific tasks and are the best services to use for the described requirements.

upvoted 4 times

AWSStudyBuddy 1 year, 7 months ago

Selected Answer: B

Although tracking configuration changes and recording API calls are not intended uses for Amazon CloudWatch, it can be utilized for monitoring and collecting data. AWS CloudTrail and AWS Config are purpose-built for these specific tasks and are the best services to use for the described requirements.

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: B

CloudWatch is a monitoring service for AWS resources and applications. CloudTrail is a web service that records API activity in your AWS account.

upvoted 2 times

Bogs123456711 1 year, 10 months ago

Selected Answer: B

CONFIG - AWS CONFIG
RECORD API CALLS - CLOUDTRAIL

upvoted 3 times

hsinchang 1 year, 10 months ago

Selected Answer: B

CloudWatch is mainly used to monitor AWS services with metrics, not recording actions inside the AWS environments. It can also monitor CloudTrail logged events.
For recording API calls it requires CloudTrail.

upvoted 2 times

james2033 1 year, 10 months ago

Selected Answer: B

Keyword "Amazon CloudWatch" is not for this case, remove C and D.

Use AWS Config first to track configuration changes, Second is AWS CloudTrail to record API calls. (Answer B, and correct answer). Answer A is reversed order of B, and not accepted.

upvoted 3 times

miki111 1 year, 11 months ago

Option B is the right answer for this.

upvoted 1 times

karloscetina007 1 year, 11 months ago

Selected Answer: B

B is the answer with no doubts

upvoted 1 times

minhpn 1 year, 11 months ago

Selected Answer: B

config => AWS config

record API calls => AWS CloudTrail

upvoted 1 times

cookieMr 1 year, 12 months ago

Selected Answer: B

To meet the requirement of tracking configuration changes on AWS resources and recording a history of API calls, a solutions architect should recommend option B: Use AWS Config to track configuration changes and AWS CloudTrail to record API calls.

Option A (using CloudTrail to track configuration changes and Config to record API calls) is incorrect because CloudTrail is specifically designed to capture API call history, while Config is designed for tracking configuration changes.

Option C (using Config to track configuration changes and CloudWatch to record API calls) is not the recommended approach. While CloudWatch can be used for monitoring and logging, it does not provide the same level of detail and compliance tracking as CloudTrail for recording API calls.

Option D (using CloudTrail to track configuration changes and CloudWatch to record API calls) is not the optimal choice because CloudTrail is the appropriate service for tracking configuration changes, while CloudWatch is not specifically designed for recording API call history.

upvoted 3 times

Bmarodi 2 years ago

Selected Answer: B

Option B meets requirements.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #35

Topic 1

A company is preparing to launch a public-facing web application in the AWS Cloud. The architecture consists of Amazon EC2 instances within a VPC behind an Elastic Load Balancer (ELB). A third-party service is used for the DNS. The company's solutions architect must recommend a solution to detect and protect against large-scale DDoS attacks. Which solution meets these requirements?

- A. Enable Amazon GuardDuty on the account.
- B. Enable Amazon Inspector on the EC2 instances.
- C. Enable AWS Shield and assign Amazon Route 53 to it.
- D. Enable AWS Shield Advanced and assign the ELB to it. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

ninjawr **Highly Voted** 2 years, 8 months ago

Selected Answer: D

Answer is D
C is incorrect because question says Third party DNS and route 53 is AWS proprietary
upvoted 43 times

kidomaruto 1 year, 7 months ago

Right answer, wrong explanation.
You can use Route 53 with a custom domain.. it's all about the "large-scale DDOS attack".
upvoted 22 times

BoboChow **Highly Voted** 2 years, 8 months ago

Selected Answer: D

AWS Shield Advanced provides expanded DDoS attack protection for your Amazon EC2 instances, Elastic Load Balancing load balancers, CloudFront distributions, Route 53 hosted zones, and AWS Global Accelerator standard accelerators.
upvoted 30 times

leopardh 2 years, 1 month ago

reunaron 2 years, 1 month ago

I'd agree as Shield Advanced is the only tier that can protect EC2 which is not possible in Standard.

upvoted 7 times

satyaamm Most Recent 5 months, 2 weeks ago

Selected Answer: D

Only AWS Shield Advanced can help against DDos attacks and since we are using a third-party DNS so we can't attach it to Route 53.

upvoted 3 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: D

Enabling AWS Shield Advanced provides advanced protection against large-scale DDoS attacks.

Assigning the ELB to AWS Shield Advanced ensures that traffic to the web application is monitored and protected with enhanced mitigation techniques.

The solution also offers a direct response mechanism (via the AWS DDoS Response Team) for further assistance during attacks.

upvoted 1 times

Tjazz04 6 months ago

Selected Answer: D

Ans. D - Keyword - Large-scale DDoS attacks which the AWS Shield Advanced can prevent

upvoted 1 times

zied007 9 months, 1 week ago

Selected Answer: D

Answer is D

upvoted 1 times

PaulGa 9 months, 2 weeks ago

Selected Answer: D

Ans D. Shield (Advanced) is built for DDoS and can interface to ELB

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

A: GuardDuty is not for this, mostly for account monitoring for suspicious activity

B: Inspector is for OS vulnerabilities

C: Shield with R53 is not going to protect against DDoS

D: Shield Advanced is build for DDoS protection

upvoted 7 times

awsgeek75 1 year, 5 months ago

Forgot to mention, C won't work because a 3rd party DNS is used and R53 is not part of the setup

upvoted 4 times

awsgeek75 1 year, 5 months ago

Prevent large scale DDOS attack = AWS Shield Advanced

upvoted 2 times

A_jaa 1 year, 5 months ago

Selected Answer: D

Answer-D

upvoted 1 times

djgodzilla 1 year, 6 months ago

Selected Answer: D

- In addition to the network and transport layer protections that come with Standard, Shield Advanced provides additional

detection and mitigation against large and sophisticated DDoS attacks, near real-time visibility into attacks, and integration with AWS WAF, a web application firewall.

<https://aws.amazon.com/shield/features/#:~:text=In%20addition%20to%20the%20network,WAF%2C%20a%20web%20application%20firewall.>

upvoted 2 times

OmegaLambda7XL9 1 year, 6 months ago

This one got me to be honest

upvoted 3 times

Ruffyt 1 year, 7 months ago

Option A is incorrect because Amazon GuardDuty is a threat detection service that focuses on identifying malicious activity and unauthorized behavior within AWS accounts. While it is useful for detecting various security threats, it does not specifically address large-scale DDoS attacks.

Option B is also incorrect because Amazon Inspector is a vulnerability assessment service that helps identify security issues and vulnerabilities within EC2. It does not directly protect against DDoS attacks.

Option C is not the optimal choice because AWS Shield provides basic DDoS protection for resources such as Elastic IP addresses, CloudFront, and Route53 hosted zones. However, it

upvoted 4 times

Ruffyt 1 year, 7 months ago

does not provide the advanced capabilities and assistance offered by AWS Shield Advanced, which is better suited for protecting against large-scale DDoS attacks.

Therefore, option D with AWS Shield Advanced and assigning the ELB to it is the recommended solution to detect and protect against large-scale DDoS attacks in the architecture described.

upvoted 3 times

Abitek007 1 year, 8 months ago

D, but can be tricky, the third party negates Route53

upvoted 1 times

Ak9kumar 1 year, 8 months ago

Answer D. Learn section on AWS Advanced Shield on [aws.amazon.com](https://aws.amazon.com/shield/advanced/) to help you understand this. It helped me.

upvoted 2 times

ishant101 1 year, 9 months ago

answer is D

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: D

DDos = AWS Shield

upvoted 2 times


Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions**Question #36**

Topic 1

A company is building an application in the AWS Cloud. The application will store data in Amazon S3 buckets in two AWS Regions. The company must use an AWS Key Management Service (AWS KMS) customer managed key to encrypt all data that is stored in the S3 buckets. The data in both S3 buckets must be encrypted and decrypted with the same KMS key. The data and the key must be stored in each of the two Regions.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Configure replication between the S3 buckets.
- B. Create a customer managed multi-Region KMS key. Create an S3 bucket in each Region. Configure replication between the S3 buckets. Configure the application to use the KMS key with client-side encryption. **Most Voted**
- C. Create a customer managed KMS key and an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Configure replication between the S3 buckets.
- D. Create a customer managed KMS key and an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with AWS KMS keys (SSE-KMS). Configure replication between the S3 buckets.

Correct Answer: B*Community vote distribution***Comments****pooppants** **Highly Voted** 2 years, 8 months ago**Selected Answer: B**

KMS Multi-region keys are required <https://docs.aws.amazon.com/kms/latest/developerguide/multi-region-keys-overview.html>
upvoted 71 times

hypnozz 2 years ago

The answer is C, because "Server-side encryption with Amazon S3 managed keys (SSE-S3) is the base level of encryption configuration for every bucket in Amazon S3. If you want to use a different type of default encryption, you can also specify server-side encryption with AWS Key Management Service (AWS KMS) keys (SSE-KMS) or customer-provided keys (SSE-C)"

By using SSE-KMS, you can encrypt the data stored in the S3 buckets with a customer managed KMS key. This ensures that the data is protected and allows you to have control over the encryption key. By creating an S3 bucket in each Region and

configuring replication between them, you can have data and key redundancy in both Regions.

upvoted 5 times

Clouddon 1 year, 10 months ago

Option B, AWS KMS supports multi-Region keys, which are AWS KMS keys in different AWS Regions that can be used interchangeably – as though you had the same key in multiple Regions. Each set of related multi-Region keys has the same key material and key ID, so you can encrypt data in one AWS Region and decrypt it in a different AWS Region without re-encrypting or making a cross-Region call to AWS KMS.

You can use multi-Region keys with client-side encryption libraries, such as the AWS Encryption SDK, the DynamoDB Encryption Client, and Amazon S3 client-side encryption. For an example of using multi-Region keys with Amazon DynamoDB global tables and the DynamoDB Encryption Client, see [Encrypt global data client-side with AWS KMS multi-Region keys in the AWS Security Blog](https://docs.aws.amazon.com/kms/latest/developerguide/multi-region-keys-overview.html). <https://docs.aws.amazon.com/kms/latest/developerguide/multi-region-keys-overview.html>

upvoted 7 times

aatikah 6 months ago

Does NOT meet the requirement of using a customer managed KMS key

upvoted 1 times

magazz 2 years, 6 months ago

Amazon S3 cross-region replication decrypts and re-encrypts data under a KMS key in the destination Region, even when replicating objects protected by a multi-Region key. So stating that Amazon S3 cross-region replication decrypts and re-encrypts data under a KMS key in the destination Region, even when replicating objects protected by a multi-Region key is required is incorrect

upvoted 5 times

thanhvx1 2 years, 2 months ago

Option B involves configuring the application to use client-side encryption, which can increase the operational overhead of managing and securing the keys.

upvoted 2 times

TuLe 2 years, 6 months ago

@magazz: it's not true then. Based on the document from AWS

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/replication-config-for-kms-objects.html> , we will need to setup the replication rule with destination KMS. In order to have the key available in more than 2, then multi-region key should be required. But I'm still not favor option B - we can use server-side when why wasting effort to do client side encryption.

upvoted 2 times

TuLe 2 years, 6 months ago

I would say it's true... Not sure the previous one say "not true" :D.

upvoted 1 times

JayBee65 2 years, 6 months ago

It's not clear what you are saying. Are you saying that B is correct or D is correct?

upvoted 2 times

karbob 2 years, 5 months ago

:D => is smile i thought

upvoted 2 times

Edwars 1 year ago

I'd say D because multi-region keys can be used with server side encryption as well.

"Multi-Region keys are supported in the AWS KMS console, the AWS KMS API, the AWS Encryption SDK, Amazon DynamoDB Encryption Client, and Amazon S3 Encryption Client. AWS services also let you configure multi-Region keys for server-side encryption in case you want the same key to protect data that needs both server-side and client-side encryption."

<https://aws.amazon.com/blogs/security/encrypt-global-data-client-side-with-aws-kms-multi-region-keys/>

upvoted 3 times

Instantqueue 1 year, 8 months ago

It's not correct because the question asks for server side encryption, not client side (before the objects reach the bucket).
upvoted 8 times

babayomi 9 months, 1 week ago

Note, The question did not ask for server side encryption, It stated Customer manages encryption. So reason why Customer side encryption makes the solution much easier and applicable. Hence B as an answer
upvoted 1 times

kelmryan1 1 year, 1 month ago

It says encrypt all data and the data originates from the application. Making it B
upvoted 1 times

KJa Highly Voted 2 years, 8 months ago

Selected Answer: D

Cannot be A - question says customer managed key
Cannot B - client side encryption is operational overhead
Cannot C -as it says SSE-S3 instead of customer managed
so the answer is D though it required one time setup of keys
upvoted 69 times

th3cookie 2 years, 6 months ago

How does client side encryption increase OPERATIONAL overhead? Do you think every connected client is sitting there with gpg cli, decrypting/encrypting every packet that comes in/out? No, it's done via SDK ->
<https://docs.aws.amazon.com/encryption-sdk/latest/developer-guide/introduction.html>

The correct answer is B because that's the only way to actually get the same key across multiple regions with minimal operational overhead
upvoted 14 times

kakka22 2 years, 2 months ago

"The data in both S3 buckets must be encrypted and decrypted with the same KMS key"
Client side encryption means that key is generated in from the client without storing that in the KMS...
upvoted 7 times

mattlai 2 years, 8 months ago

fun joke, if u dont do encryption on client side, where else could it be?
upvoted 1 times

Newptone 2 years, 7 months ago

It could be server side. For client side, the application need to finish the encryption and decryption by itself. So S3 object encryption on the server side is less operational overhead.
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingClientSideEncryption.html>

But for option B, the major issue is if you create KMS keys in 2 regions, they can not be the same.
upvoted 7 times

Newptone 2 years, 7 months ago

Sorry for the typo, I mean option D.
upvoted 2 times

Clouddon 1 year, 10 months ago

Kindly point at where server-side encryption support multi-region. It is only mention on the aws blog that client-side support multi-region.
upvoted 2 times

BoboChow 2 years, 8 months ago

The data in both S3 buckets must be encrypted and decrypted with the same KMS key.
AWS KMS supports multi-Region keys, which are AWS KMS keys in different AWS Regions that can be used interchangeably – so though you had the same key in multiple Regions

as though you had the same key in multiple regions.

"as though" means it's different.

So I agree with B

upvoted 15 times

BoboChow 2 years, 8 months ago

key change across regions unless you use multi-Region keys

upvoted 2 times

pentium75 1 year, 5 months ago

B includes replicating the data in the S3 buckets, which is not mentioned anywhere in the stem. It says that you need to store data in two buckets, not that you need to replicate content between buckets.

upvoted 1 times

Drew3000 1 year, 3 months ago

All the choices involve replication between the buckets.

upvoted 1 times

Yak_Yeti Most Recent 2 months, 2 weeks ago

Selected Answer: B

D says "Create a customer managed KMS key and an S3 bucket in each Region." If there are keys created in both regions, then the keys are not the same. If it said create a multi region key and synch it, then I would pick D.

upvoted 1 times

SleeplessHossam 3 months ago

Selected Answer: B

I was confused which to choose, B or D, but someone in discussion mentioned really important point, D suggests creating two keys and two regions, which they can't be the same Key, so D of course is not the correct answer.

upvoted 1 times

Anastesas 3 months, 2 weeks ago

Selected Answer: D

Given the requirements to encrypt and decrypt data stored in S3 buckets in two AWS Regions using the same KMS key with minimal operational overhead, server-side encryption with AWS KMS (SSE-KMS) is the most suitable solution. It provides simplicity, seamless integration, centralized key management, and optimal performance.

Therefore: Option D: Create a customer managed KMS key and an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with AWS KMS keys (SSE-KMS). Configure replication between the S3 buckets.

upvoted 1 times

iwantcertificates 3 months, 2 weeks ago

Selected Answer: D

I think this question would be invalidated in an actual test. Both B and D has problems. B uses client-side encryption, which adds complexity to the operation. Additionally, KMS by default uses server side encryption, SSE-KMS. But the option D also has a imprecision, because you are creating two KMS keys instead of just one with multi-region

upvoted 1 times

sammo08 3 months, 3 weeks ago

Selected Answer: B

Key Reason: Separate KMS Keys in Each Region

In Option D, you are creating a separate AWS KMS key in each Region, meaning the encryption keys are distinct.

When replicating data between Regions, AWS does not allow direct decryption of objects encrypted with an SSE-KMS key in one Region using a different KMS key in another Region.

KMS keys are tied to a specific Region unless you use a Multi-Region key.

Because of this, the replicated objects cannot be automatically decrypted in the destination Region without additional configuration, increasing operational overhead.

Why Option B is Better

Option B uses a Multi-Region KMS Key that allows encryption and decryption with the same key across both Regions.

Less operational complexity, as the application does not have to handle different keys in each Region.

Seamless replication and decryption of data across Regions.

upvoted 1 times

Clpsz 4 months, 1 week ago

Selected Answer: D

B: client side encryption is not correct

upvoted 1 times

AshishDhole 4 months, 1 week ago

Selected Answer: B

he data in both S3 buckets must be encrypted and decrypted with the same KMS key. The data and the key must be stored in each of the two Regions --> Single Key for both region. So multi region key will be the answer

upvoted 1 times

Krenil 4 months, 1 week ago

Selected Answer: D

Correct Ans: D

Option B is not the best solution because it involves client-side encryption, which adds unnecessary complexity and operational overhead. Here's why:

Client-side encryption means that the application itself must handle the encryption and decryption of data before it is uploaded to S3 and after it is downloaded. This requires additional development work, key management, and the management of encryption and decryption logic in the application.

In contrast, SSE-KMS (option D) allows Amazon S3 to manage the encryption and decryption process automatically on the server side, which reduces operational overhead and simplifies the implementation.

upvoted 1 times

Tika57 4 months, 1 week ago

Selected Answer: D

B and D are in short list (using customer managed key and KMS) but B requires that each client encrypts the data before sending it, whereas the D is totally transparent so with "LEAST overhead"

upvoted 1 times

dariar 4 months, 2 weeks ago

Selected Answer: D

For me the answer is D.

It could also be the B-answer, but when it asks "Configure the application", it sounds like the "operational overhead", so it's eliminatory for me

upvoted 1 times

AtiiF 4 months, 3 weeks ago

Selected Answer: D

=====> D

upvoted 1 times

V2910 4 months, 3 weeks ago

Selected Answer: D

Difference between B and D is server side encryption

upvoted 1 times

CloudExpert01 4 months, 3 weeks ago

Selected Answer: B

Mentioning of the "multi-Region KMS key" <https://docs.aws.amazon.com/kms/latest/developerguide/multi-region-keys-overview.html>

upvoted 1 times

aefuen1 5 months, 3 weeks ago

Selected Answer: B

It's B. it has to use the same key in both regions.

key, it has to use the same key in each region.

upvoted 2 times

pg6058 5 months, 4 weeks ago

Selected Answer: B

Multi-Region KMS keys supports client-side encryption

<https://aws.amazon.com/about-aws/whats-new/2021/06/kms-multi-region-keys/>

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #37

Topic 1

A company recently launched a variety of new workloads on Amazon EC2 instances in its AWS account. The company needs to create a strategy to access and administer the instances remotely and securely. The company needs to implement a repeatable process that works with native AWS services and follows the AWS Well-Architected Framework. Which solution will meet these requirements with the LEAST operational overhead?

A. Use the EC2 serial console to directly access the terminal interface of each instance for administration.


B. Attach the appropriate IAM role to each existing instance and new instance. Use AWS Systems Manager Session Manager to establish a remote SSH session. **Most Voted**

C. Create an administrative SSH key pair. Load the public key into each EC2 instance. Deploy a bastion host in a public subnet to provide a tunnel for administration of each instance.

D. Establish an AWS Site-to-Site VPN connection. Instruct administrators to use their local on-premises machines to connect directly to the instances by using SSH keys across the VPN tunnel.

Correct Answer: B

Community vote distribution



Comments

BoboChow **Highly Voted** 2 years, 7 months ago

Selected Answer: B

How can Session Manager benefit my organization?

Ans: No open inbound ports and no need to manage bastion hosts or SSH keys

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html>

upvoted 24 times

Nightducky 2 years, 7 months ago

Do you know what from the question is it Windows or Linux EC2. I think not so how you want to do SSH session for Windows?

Answer is C

upvoted 3 times

JayBee65 2 years, 6 months ago

Session Manager provides support for Windows, Linux, and macOS from a single tool

session manager provides support for windows, linux, and macOS from a single tool
upvoted 10 times

sohain 1 year, 10 months ago

session manager works with linux, windows, and mac too
upvoted 4 times

TienHuynh 1 year, 11 months ago

"Cross-platform support for Windows, Linux, and macOS"
<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html>

upvoted 3 times

cookieMr Highly Voted 1 year, 12 months ago

Selected Answer: B

Option A provides direct access to the terminal interface of each instance, but it may not be practical for administration purposes and can be cumbersome to manage, especially for multiple instances.

Option C adds operational overhead and introduces additional infrastructure that needs to be managed, monitored, and secured. It also requires SSH key management and maintenance.

Option D is complex and may not be necessary for remote administration. It also requires administrators to connect from their local on-premises machines, which adds complexity and potential security risks.

Therefore, option B is the recommended solution as it provides secure, auditable, and repeatable remote access using IAM roles and AWS Systems Manager Session Manager, with minimal operational overhead.

upvoted 9 times

Kazmin Most Recent 2 months, 2 weeks ago

Selected Answer: A

looks correct

upvoted 1 times

Rcosmos 5 months ago

Selected Answer: B

A Opção B é a solução mais simples, segura e de menor sobrecarga operacional. O uso do AWS Systems Manager Session Manager elimina a necessidade de abrir portas, gerenciar chaves SSH ou configurar infraestrutura adicional, alinhando-se ao AWS Well-Architected Framework.

upvoted 1 times

satyaamm 5 months, 2 weeks ago

Selected Answer: B

Creating an IAM Role is the most suitable as others do not meet the criteria of less operational overhead.

upvoted 1 times

safa_123 7 months, 1 week ago

B. Create a customer managed multi-Region KMS key. Create an S3 bucket in each Region. Configure replication between the S3 buckets. Configure the application to use the KMS key with client-side encryption.

Explanation:

Multi-Region KMS Key: AWS KMS supports multi-Region keys, which can be replicated across Regions, enabling encryption and decryption in multiple Regions using the "same" KMS key (though technically it's a replica). This meets the requirement of using the same key in both Regions without additional management overhead for separate keys.

Client-side encryption: Configuring the application to use the KMS key with client-side encryption ensures that the data is encrypted before it is sent to Amazon S3, and decrypted when retrieved.

upvoted 2 times

MortisG 1 month ago

This is perfect for the previous question.

upvoted 1 times

zied007 9 months, 1 week ago

Selected Answer: B

Answer is B

upvoted 1 times

ManikRoy 1 year, 1 month ago

Selected Answer: B

Session Manager provides secure and auditable node management without the need to open inbound ports, maintain bastion hosts, or manage SSH keys.

upvoted 4 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

A: Serial console is for device direct connection to peripherals and monitor boot etc.

C: Workable solution but a lot of overhead

D: Too much overhead for everyone

B: Managed product for this purpose so least overhead.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html>

upvoted 2 times

A_jaa 1 year, 5 months ago

Selected Answer: B

Answer-B

upvoted 1 times

AWSSStudyBuddy 1 year, 7 months ago

I go with option B. Here's why--- IAM Roles: Without SSH keys or shared passwords, securely provide access to EC2 instances and AWS services.

upvoted 4 times

AWSSStudyBuddy 1 year, 7 months ago

Without requiring direct SSH connection, securely access and control EC2 instances with AWS Systems Manager Session Manager.

Least Operational Overhead: An effective and fully managed method of managing instances.

Well-Architected Framework: Complies with performance, security, and reliability best practices from AWS.

Cons of alternative options:

Option A: The automation and flexibility required for secure administration at scale are not provided by using the EC2 serial terminal directly.

Option C: There is more operational overhead and complexity when a bastion host is deployed.

Option D: For secure instance administration, setting up an AWS Site-to-Site VPN connection is too difficult and not the optimal approach.

In conclusion, Option B is suggested as the best option given the given circumstances.

upvoted 5 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

This solution meets all of the requirements with the LEAST operational overhead. It is repeatable, uses native AWS services, and follows the AWS Well-Architected Framework.

Repeatable: The process of attaching an IAM role to an EC2 instance and using Systems Manager Session Manager to establish a remote SSH session is repeatable. This can be easily automated, so that new instances can be provisioned and administrators can connect to them securely without any manual intervention.

upvoted 3 times

TanishKishore 1 year, 10 months ago

tanqnpkemei 1 year, 10 months ago

Selected Answer: B

With AWS Systems Manager Session Manager, you can manage your Amazon Elastic Compute Cloud (Amazon EC2) instances, edge devices, on-premises servers, and virtual machines (VMs). You can use either an interactive one-click browser-based shell or the AWS Command Line Interface (AWS CLI). It provides secure and auditable node management without the need to open inbound ports, maintain bastion hosts, or manage SSH keys.

<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html#:~:text=RSS-,Session%20Manager,-is%20a%20fully>

upvoted 3 times

james2033 1 year, 10 months ago

Selected Answer: B

Keyword "access and administer the instances remotely and securely" See "AWS Systems Manager Session Manager at "<https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html> .

upvoted 2 times

miki111 1 year, 11 months ago

Option B is the right answer for this.

upvoted 1 times

TienHuynh 1 year, 11 months ago

Selected Answer: B

- +Centralized access control to managed nodes using IAM policies
- +No open inbound ports and no need to manage bastion hosts or SSH keys
- +Cross-platform support for Windows, Linux, and macOS

upvoted 2 times

Bmarodi 2 years ago

Selected Answer: B

The choice for me is the option B.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #38

Topic 1

A company is hosting a static website on Amazon S3 and is using Amazon Route 53 for DNS. The website is experiencing increased demand from around the world. The company must decrease latency for users who access the website. Which solution meets these requirements MOST cost-effectively?

- A. Replicate the S3 bucket that contains the website to all AWS Regions. Add Route 53 geolocation routing entries.
- B. Provision accelerators in AWS Global Accelerator. Associate the supplied IP addresses with the S3 bucket. Edit the Route 53 entries to point to the IP addresses of the accelerators.
- C. Add an Amazon CloudFront distribution in front of the S3 bucket. Edit the Route 53 entries to point to the CloudFront distribution. **Most Voted**
- D. Enable S3 Transfer Acceleration on the bucket. Edit the Route 53 entries to point to the new endpoint.

Correct Answer: C

Community vote distribution



Comments

cookieMr **Highly Voted** 1 year, 12 months ago

Selected Answer: C

Option A (replicating the S3 bucket to all AWS Regions) can be costly and complex, requiring replication of data across multiple Regions and managing synchronization. It may not provide a significant latency improvement compared to the CloudFront solution.

Option B (provisioning accelerators in AWS Global Accelerator) can be more expensive as it adds an extra layer of infrastructure (accelerators) and requires associating IP addresses with the S3 bucket. CloudFront already includes global edge locations and provides similar acceleration capabilities.

Option D (enabling S3 Transfer Acceleration) can help improve upload speed to the S3 bucket but may not have a significant impact on reducing latency for website visitors.

Therefore, option C is the most cost-effective solution as it leverages CloudFront's caching and global distribution capabilities to decrease latency and improve website performance.

upvoted 40 times

TariqKipkemei **Highly Voted** 1 year, 10 months ago

Keywords:

Keywords:

Global, Reduce latency, S3, Static Website, Cost effective = Amazon CloudFront

upvoted 5 times

AshishDhole Most Recent 5 months ago

Selected Answer: C

Cost effective solution + Static Website = Cloud Front + S3

upvoted 2 times

satyaamm 5 months, 2 weeks ago

Selected Answer: C

CloudFront helps reduce latency most cost-effectively.

upvoted 1 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: C

Global Content Delivery:

CloudFront is a content delivery network (CDN) that caches content at edge locations around the world, reducing latency by serving content from the nearest edge location to the user.

Seamless Integration:

CloudFront integrates easily with S3 buckets. It fetches and caches the static content while reducing the load on the origin S3 bucket.

Cost-Effectiveness:

CloudFront charges based on data transfer and requests, which is often more economical than replicating S3 buckets across Regions or using other solutions that require additional infrastructure.

Performance and Scalability:

CloudFront handles spikes in traffic without additional configuration or cost overhead for provisioning and managing infrastructure.

upvoted 1 times

zied007 9 months, 1 week ago

Selected Answer: C

Answer is C

upvoted 1 times

hb0011 1 year ago

Some of these questions seem too easy for SAA. These seem like Cloud Practitioner questions...

upvoted 2 times

TilTil 1 year, 2 months ago

Cloudfront is a lovely and affordable CDN for static content.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

S3 static website so CloudFront is the best CDN solution for low cost and low latency

A: Very expensive way of doing this

B: Makes no sense

D: Transfer Acc is for upload boosting

C: CloudFront literally solves this problem

upvoted 2 times

A_jaa 1 year, 5 months ago

Selected Answer: C

Answer-C

upvoted 1 times

Ruffyt 1 year, 7 months ago

Option A (replicating the S3 bucket to all AWS Regions) can be costly and complex, requiring replication of data across multiple Regions and managing synchronization. It may not provide a significant latency improvement compared to the CloudFront solution.

Option B (provisioning accelerators in AWS Global Accelerator) can be more expensive as it adds an extra layer of infrastructure (accelerators) and requires associating IP addresses with the S3 bucket. CloudFront already includes global edge locations and provides similar acceleration capabilities.

Option D (enabling S3 Transfer Acceleration) can help improve upload speed to the S3 bucket but may not have a significant impact on reducing latency for website visitors.

Therefore, option C is the most cost-effective solution as it leverages CloudFront's caching and global distribution capabilities to decrease latency and improve website performance.

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

Amazon CloudFront is a content delivery network (CDN) service that distributes content globally to reduce latency. By setting up a CloudFront distribution in front of the S3 bucket hosting the static website, you can take advantage of its edge locations around the world to deliver content from the nearest location to the users, reducing the latency they experience.

CloudFront automatically caches and replicates content to its edge locations, resulting in faster delivery and lower latency for users worldwide. This solution is highly effective in optimizing performance while keeping costs under control because CloudFront charges are based on actual data transfer and requests, and the pay-as-you-go pricing model ensures that you only pay for what you use.

upvoted 5 times

james2033 1 year, 10 months ago

Selected Answer: C

Keyword "Amazon CloudFront" (C).

upvoted 1 times

miki111 1 year, 11 months ago

Option C is the right answer for this.

upvoted 1 times

miki111 1 year, 11 months ago

Option C is the right answer for this.

upvoted 1 times

TienHuynh 1 year, 11 months ago

Selected Answer: C

key words:

- around the world
- decrease latency
- most cost-effective

answer is C

upvoted 2 times

cheese929 2 years, 2 months ago

Selected Answer: C

C is the most cost effective.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #39


Topic 1

A company maintains a searchable repository of items on its website. The data is stored in an Amazon RDS for MySQL database table that contains more than 10 million rows. The database has 2 TB of General Purpose SSD storage. There are millions of updates against this data every day through the company's website. The company has noticed that some insert operations are taking 10 seconds or longer. The company has determined that the database storage performance is the problem. Which solution addresses this performance issue?

- A. Change the storage type to Provisioned IOPS SSD. **Most Voted**
- B. Change the DB instance to a memory optimized instance class.
- C. Change the DB instance to a burstable performance instance class.
- D. Enable Multi-AZ RDS read replicas with MySQL native asynchronous replication.

Correct Answer: A

Community vote distribution



Comments

pazabal **Highly Voted** 2 years, 5 months ago

Selected Answer: A

A: Made for high levels of I/O opps for consistent, predictable performance.
 B: Can improve performance of insert opps, but it's a storage performance rather than processing power problem
 C: for moderate CPU usage
 D: for scale read-only replicas and doesn't improve performance of insert opps on the primary DB instance
 upvoted 42 times

cookieMr **Highly Voted** 1 year, 12 months ago

Selected Answer: A

Option B (changing the DB instance to a memory optimized instance class) focuses on improving memory capacity but may not directly address the storage performance issue.

Option C (changing the DB instance to a burstable performance instance class) is suitable for workloads with varying usage patterns and burstable performance needs, but it may not provide consistent and predictable performance for heavy write workloads.

Option D (enabling Multi-AZ RDS read replicas with MySQL native asynchronous replication) is a solution for high availability and read scaling but does not directly address the storage performance issue.

Therefore, option A is the most appropriate solution to address the performance issue by leveraging Provisioned IOPS SSD storage type, which provides consistent and predictable I/O performance for the Amazon RDS for MySQL database.

upvoted 27 times

JA2018 Most Recent 7 months, 1 week ago

Selected Answer: A

Hint in the stem: DB insert operations taking 10 seconds or longer => current General Purpose SSD storage cannot providing sufficient IOPS required for the workload. Provisioned IOPS SSD can significantly improve storage performance and reduce latency for DB operations.

Provisioned IOPS SSD: Designed to deliver fast, predictable, and consistent I/O performance. This type of performance is very crucial for databases with high transaction rates and frequent updates. It allows users to provision a specific level of IOPS to meet the performance requirements of their specific workload.

upvoted 1 times

PaulGa 9 months, 1 week ago

Selected Answer: A

Ans A, as per stephyfresh13:

"Provisioned IOPS SSD: This storage type is designed to deliver fast, predictable, and consistent I/O performance, which is crucial for databases with high transaction rates and frequent updates. It allows you to provision a specific level of IOPS to meet the performance requirements of your workload."

upvoted 3 times

stephyfresh13 11 months, 2 weeks ago

A. Change the storage type to Provisioned IOPS SSD.

Here's why:

Provisioned IOPS SSD: This storage type is designed to deliver fast, predictable, and consistent I/O performance, which is crucial for databases with high transaction rates and frequent updates. It allows you to provision a specific level of IOPS to meet the performance requirements of your workload.

Current Issue: The problem with insert operations taking 10 seconds or longer indicates that the current General Purpose SSD storage is not providing sufficient IOPS for the workload. Provisioned IOPS SSD can significantly improve storage performance and reduce latency for database operations.

upvoted 4 times

sudohogan 1 year, 1 month ago

B is the answer, if the company decided that storage is the problem then an IOPS SSD (storage) is the solution and not a memory optimiser.

upvoted 1 times

TiITi 1 year, 2 months ago

Selected Answer: B

A. IOPS is about increasing the number of input connections so you can handle more requests. Which may not be the issue.
B. Having a memory optimized class provides more RAM to execute the queries which take upto 10 secs to complete. More RAM means they can execute faster.

C and D are distractors. They deal with high availability and timely scalability which are not issues here.

upvoted 2 times

MortisG 1 month ago

It says insert operations are taking up to 10 secs. This is clearly I/O so optimizing memory won't help. Correct answer is A.

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

Database storage is issue so

BD is irrelevant

C is for performance boost (CPU) which won't help with storage issues

A Fix the storage issue
upvoted 2 times

A_jaa 1 year, 5 months ago

Selected Answer: A

Answer-A
upvoted 1 times

bujuman 1 year, 5 months ago

Selected Answer: B

Do not misconsider "database storage performance is the problem". I beleive the correct asnwer is B Due too the fact that Mmemory Optimized EC2 instance family is designed for DB servers perf.
upvoted 1 times

pentium75 1 year, 5 months ago

But the stem clearly says that storage performance, NOT memory performance, is the problem. More memory won't increase storage performance.
upvoted 6 times

aptx4869 1 year, 7 months ago

Selected Answer: A

A is correct answer because it is talking about storage and transaction speed is slow due to it, should change to iops storage instead.
upvoted 2 times

Ruffyt 1 year, 7 months ago

A: Made for high levels of I/O opps for consistent, predictable performance.
B: Can improve performance of insert opps, but it's a storage performance rather than processing power problem
C: for moderate CPU usage
D: for scale read-only replicas and doesn't improve performance of insert opps on the primary DB instance
upvoted 1 times

Rhydian25 1 year ago

Copy-paste from pazabal's answer...
upvoted 1 times

AWSSStudyBuddy 1 year, 7 months ago

I go with option A. Using Amazon Provisioned IOPS (PIOPS) SSD storage is the best way to solve the performance issue of insert operations taking 10 seconds or longer on an Amazon RDS for MySQL database table with more than 10 million rows and 2 TB of General Purpose SSD storage.

A high-performance storage solution with reliable throughput and minimal latency is PIOPS SSD storage. Workloads like insert operations, which demand high I/O performance, are ideally suited for it.
upvoted 3 times

tom_cruise 1 year, 8 months ago

Selected Answer: A

Key: database storage performance is the problem.
upvoted 1 times

awsleffe 1 year, 8 months ago

Selected Answer: A

Option A is answer - A. Change the storage type to Provisioned IOPS SSD.
The company's issue is related to storage performance, specifically with insert operations. This suggests that the I/O operations are the bottleneck.

Provisioned IOPS SSD storage type is designed to handle the kind of workload the company is experiencing and should help improve the performance of insert operations.
upvoted 3 times

upvoted 3 times

awashenko 1 year, 8 months ago

Selected Answer: A

"The company has determined that the database storage performance is the problem."

This is the key statement in the question. Otherwise I would have selected B but this statement here makes A correct.

upvoted 2 times

David_Ang 1 year, 8 months ago

Selected Answer: A

yeah "A" is correct is the most suitable option for this scenario, because you need to improve the speed of the reading and writing of the storage system.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #40

Topic 1

A company has thousands of edge devices that collectively generate 1 TB of status alerts each day. Each alert is approximately 2 KB in size. A solutions architect needs to implement a solution to ingest and store the alerts for future analysis.

The company wants a highly available solution. However, the company needs to minimize costs and does not want to manage additional infrastructure. Additionally, the company wants to keep 14 days of data available for immediate analysis and archive any data older than 14 days.

What is the MOST operationally efficient solution that meets these requirements?

A. Create an Amazon Kinesis Data Firehose delivery stream to ingest the alerts. Configure the Kinesis Data Firehose stream to deliver the alerts to an Amazon S3 bucket. Set up an S3 Lifecycle configuration to transition data to Amazon S3 Glacier after 14 days. **Most Voted**

B. Launch Amazon EC2 instances across two Availability Zones and place them behind an Elastic Load Balancer to ingest the alerts. Create a script on the EC2 instances that will store the alerts in an Amazon S3 bucket. Set up an S3 Lifecycle configuration to transition data to Amazon S3 Glacier after 14 days.

C. Create an Amazon Kinesis Data Firehose delivery stream to ingest the alerts. Configure the Kinesis Data Firehose stream to deliver the alerts to an Amazon OpenSearch Service (Amazon Elasticsearch Service) cluster. Set up the Amazon OpenSearch Service (Amazon Elasticsearch Service) cluster to take manual snapshots every day and delete data from the cluster that is older than 14 days.

D. Create an Amazon Simple Queue Service (Amazon SQS) standard queue to ingest the alerts, and set the message retention period to 14 days. Configure consumers to poll the SQS queue, check the age of the message, and analyze the message data as needed. If the message is 14 days old, the consumer should copy the message to an Amazon S3 bucket and delete the message from the SQS queue.

Correct Answer: A*Community vote distribution*

A (87%)

D (13%)

Comments**Sinaneos** **Highly Voted** 2 years, 2 months ago**Selected Answer: A**

Definitely A, it's the most operationally efficient compared to D, which requires a lot of code and infrastructure to maintain. A is mostly managed (firehose is fully managed and S3 lifecycles are also managed)

upvoted 41 times

Kelvin_ke 2 years ago

what about the 30 days minimum requirement to transition to S3 glacier?

upvoted 8 times

caffee 1 year, 8 months ago

This constraint is related to moving from Standard to IA/IA-One Zone only. Nothing to do with Glacier

upvoted 5 times

studis 1 year, 12 months ago

You can directly migrate from S3 standard to glacier without waiting

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 5 times

jerryI 2 months, 3 weeks ago

the 30 day requirement seem only apply to IA but the article did not mention any limitation on glacier

upvoted 1 times

ErnShm 1 year, 6 months ago

the current article doesn't enable the current option, minimum days are 30

upvoted 1 times

Suvam90 1 year, 3 months ago

No , It's not correct , We can change the storage class in day 0 also using lifecycle policy , I implemented in my project, 30 days is just an example.

upvoted 5 times

Abrar2022 1 year, 6 months ago

GLACIER IS 7 DAYS REQUIREMENT NOT 30

upvoted 7 times

123jh10 Highly Voted 2 years, 1 month ago

Selected Answer: A

Only A makes sense operationally.

If you think D, just consider what is needed to move the message from SQS to S3... you are polling daily 14 TB to take out 1 TB... that's no operationally efficient at all.

upvoted 21 times

satyaamm Most Recent 5 months, 2 weeks ago

Selected Answer: A

Using Kinesis Data Firehouse with S3 Life Configuration policies is the most suitable option as we are looking for less operational overhead.

upvoted 1 times

Parul25 10 months, 3 weeks ago

I understand A is the most operationally efficient option of all but I can't wrap my head around the fact that objects must have a minimum of 30 days before they can transition or expire from Amazon S3. This means that for the first 30 days after an item is created, you cannot transition or remove it. So, how option A can be the best fit?

upvoted 1 times

cheroh_tots 10 months ago

The same 30-day minimum applies when you specify a transition from S3 Standard-IA storage to S3 One Zone-IA.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 3 times

krx59456zslszcom 5 months, 1 week ago

But the docs are also mentioning a table with minimum storage durations beyond the S3 Standard-IA
"[...]for a list of minimum storage duration for all storage class, see Comparing the Amazon S3 storage classes."
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/storage-class-intro.html#sc-compare>
upvoted 1 times

awsgeek75 11 months ago

Selected Answer: A

BCD: Additional infra which company doesn't want
A: Firehose for ingestion and delivery to S3. Lifecycle for managing archive. Fully managed and operationally easy solution
upvoted 2 times

A_jaa 11 months ago

Selected Answer: A

Answer-A
upvoted 1 times

jjcode 1 year ago

so many words...
upvoted 4 times

OmegaLambda7XL9 1 year ago

That was an easy A. Kinesis Firehose can load data directly to S3 which makes it the most operationally efficient
upvoted 2 times

Ruffyt 1 year, 1 month ago

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>
upvoted 2 times

tom_cruise 1 year, 2 months ago

Selected Answer: A

Key: MOST operationally efficient solution
upvoted 2 times

David_Ang 1 year, 2 months ago

Selected Answer: A

"A" is simply correct because kinesis firehouse is made for this, SQS standard is not going to support 500 million alerts 2KB each (1 TB) this service is made for requests that are lighter.
upvoted 2 times

Ak9kumar 1 year, 2 months ago

I picked A. Appeared to be right answer.
upvoted 1 times

chandu7024 1 year, 2 months ago

Should be A
upvoted 1 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: A

The MOST operationally efficient option is A.
upvoted 1 times

james2033 1 year, 4 months ago

Selected Answer: A

Keyword "Amazon S3 Glacier" (A).

negotiate Amazon EC2 cluster capacity.
upvoted 1 times

miki111 1 year, 4 months ago

Option A is the right answer for this.
upvoted 1 times

cookieMr 1 year, 5 months ago

Selected Answer: A

B suggests launching EC2 instances to ingest and store the alerts, which introduces additional infrastructure management overhead and may not be as cost-effective and scalable as using managed services like Kinesis Data Firehose and S3.

C involves delivering the alerts to an Amazon OpenSearch Service cluster and manually managing snapshots and data deletion. This introduces additional complexity and manual overhead compared to the simpler solution of using Kinesis Data Firehose and S3.

D suggests using SQS to ingest the alerts, but it does not provide the same level of data persistence and durability as storing the alerts directly in S3. Additionally, it requires manual processing and copying of messages to S3, which adds operational complexity.

Therefore, A provides the most operationally efficient solution that meets the company's requirements by leveraging Kinesis Data Firehose to ingest the alerts, storing them in an S3 bucket, and using an S3 Lifecycle configuration to transition data to S3 Glacier for long-term archival, all without the need for managing additional infrastructure.

upvoted 11 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions**Question #41****Topic 1**

A company's application integrates with multiple software-as-a-service (SaaS) sources for data collection. The company runs Amazon EC2 instances to receive the data and to upload the data to an Amazon S3 bucket for analysis. The same EC2 instance that receives and uploads the data also sends a notification to the user when an upload is complete. The company has noticed slow application performance and wants to improve the performance as much as possible.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Auto Scaling group so that EC2 instances can scale out. Configure an S3 event notification to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete.


B. Create an Amazon AppFlow flow to transfer data between each SaaS source and the S3 bucket. Configure an S3 event notification to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete. **Most Voted**

C. Create an Amazon EventBridge (Amazon CloudWatch Events) rule for each SaaS source to send output data. Configure the S3 bucket as the rule's target. Create a second EventBridge (Cloud Watch Events) rule to send events when the upload to the S3 bucket is complete. Configure an Amazon Simple Notification Service (Amazon SNS) topic as the second rule's target.

D. Create a Docker container to use instead of an EC2 instance. Host the containerized application on Amazon Elastic Container Service (Amazon ECS). Configure Amazon CloudWatch Container Insights to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete.

Correct Answer: B

Community vote distribution

**Comments**

jdr75 **Highly Voted** 2 years, 2 months ago

Selected Answer: A

It says "LEAST operational overhead" (ie do it in a way it's the less work for me).

If you know a little Amazon AppFlow (see the some videos) you'll see you'll need time to configure and test it, and at the end cope with the errors during the extraction and load the info to the target.

The customer in the example ALREADY has some EC2 that do the work, the only problem is the performance, that WILL be improved scaling out and adding a queue (SNS) to decouple the work of notify the user.

The operational load of doing this is LESS that configuring AppFlow.

upvoted 47 times

nhaastrup 5 months, 3 weeks ago

No B. if you go to the aws site(<https://aws.amazon.com/appflow/>) -AppFlow is a fully managed integration service that helps you securely transfer data between software as a service (SaaS) applications such as Salesforce, SAP, Google Analytics, Facebook Ads, and ServiceNow, and AWS services such as Amazon Simple Storage Service (S3) and Amazon Redshift in just a "few clicks". remember the question is asking for "the LEAST operational overhead", While Auto Scaling improves performance, it still involves managing EC2 instances, which adds operational overhead compared to a managed service like AppFlow.

upvoted 6 times

the_mellie 1 year ago

Agreed, I hesitated a long time exactly for this reason. However, the requirement is to improve the performance as much as possible. I cannot guarantee the same performance with EC2 than with managed services. That's why I eventually voted B.

upvoted 2 times

pentium75 1 year, 5 months ago

"Operational overhead" refers to operation of the solution once it's deployed, it's not about setting it up. EC2 instances that retrieve data from A and write it to B are nonsense, that's what cloud services are meant for.

upvoted 14 times

awsgeek75 1 year, 5 months ago

With AWS, a managed service is "less operational overhead" regardless of the complexity of the setup. AppFlow management is less of a headache when compared to EC2 management so A cannot be correct. EC2 has a setup overhead of OS/Application/Code hooks, security etc. continuous patching/upgrading seems like more than what you'll need to do with SaaS.

B is the correct answer.

upvoted 9 times

Six_Fingered_Jose Highly Voted 2 years, 7 months ago

Selected Answer: B

This question just screams AppFlow (SaaS integration)

<https://aws.amazon.com/appflow/>

upvoted 40 times

Six_Fingered_Jose 2 years, 7 months ago

configuring Auto-Scaling also takes time when compared to AppFlow, in AWS's words "in just a few clicks"

> Amazon AppFlow is a fully managed integration service that enables you to securely transfer data between Software-as-a-Service (SaaS) applications like Salesforce, SAP, Zendesk, Slack, and ServiceNow, and AWS services like Amazon S3 and Amazon Redshift, in just a few clicks

upvoted 23 times

Kazmin Most Recent 2 months, 2 weeks ago

Selected Answer: A

A is correct

upvoted 1 times

Itachi28 2 months, 3 weeks ago

Selected Answer: B

Option A: Auto Scaling EC2 instances

High operational overhead (still managing EC2 instances).

Not eliminating bottlenecks (EC2 instance is still the data processor).

□ Option C: EventBridge Rules

Complex event-driven setup instead of direct integration.

Requires custom event rules for each SaaS source, increasing operational burden.

□ Option D: ECS with CloudWatch Events

Migrating to ECS introduces complexity and doesn't directly improve transfer speeds.

Still processes data manually, adding an unnecessary layer

still processes data manually, adding an unnecessary layer.

Amazon AppFlow (Option B) is the best choice because it is fully managed, removes EC2 dependencies, and simplifies integration while improving performance.

upvoted 1 times

Dharmarajan 4 months, 2 weeks ago

Selected Answer: B

Initially I was inclined towards A, but then I read about AppFlow, which is a specific service to collect SAAS data sources. Also with A, there will be additional maintenance activities for EC2 and application, but AppFlow will completely eliminate that work!

upvoted 1 times

Skyskilo 5 months ago

Selected Answer: A

Key words here is Least Operational Overhead so A for me

upvoted 1 times

vuongdo 5 months, 3 weeks ago

Selected Answer: B

The answer is B, because for answer A you will have to manage the EC2 infrastructure, while AppFlow is a serverless service, you will not need to manage the infrastructure

upvoted 3 times

Ilccing 6 months ago

Selected Answer: B

i support AppFlow, the original can be saas, and the target can be s3.

upvoted 2 times

f51a8bd 7 months ago

B. Cree un flujo de Amazon AppFlow para transferir datos entre cada fuente de SaaS y el depósito de S3. Configure una notificación de eventos de S3 para enviar eventos a un tema de Amazon Simple Notification Service (Amazon SNS) cuando se complete la carga al depósito de S3.

Explicación:

Amazon AppFlow permite la integración directa con múltiples fuentes de SaaS y automatiza el flujo de datos hacia Amazon S3, eliminando la necesidad de instancias de EC2 para recibir y cargar datos. Esto reduce la carga en las instancias de EC2 y simplifica el proceso de integración y recolección de datos.

Notificación de eventos de S3 y SNS : Configurar eventos en Amazon S3 permite enviar notificaciones directamente a un tema de Amazon SNS cuando los datos se han cargado, lo cual asegura que los usuarios recibirán notificaciones sin necesidad de código personalizado en EC2.

upvoted 2 times

ChymKuBoy 8 months ago

Selected Answer: B

B for sure

upvoted 1 times

rar75024 8 months, 1 week ago

direct connect maybe correct answer but we are making too many assumptions regarding direct connect network requirements.

upvoted 1 times

effiecancode 11 months, 1 week ago

Has to be appflow because of the SaaS integration

upvoted 2 times

jatric 11 months, 2 weeks ago

Selected Answer: B

First i thought it would be A but with more study found if its SaaS and need to choose between EC2 and AWS managed service one should always choose AWS managed service. So AppFlow seems more appropriate here.

upvoted 4 times

SMALLE 1 year, 4 months ago

Selected Answer: B

Amazon AppFlow is a fully managed integration service that helps you securely transfer data between software as a service (SaaS) applications such as Salesforce, SAP, Google Analytics, Facebook Ads, and ServiceNow, and AWS services such as Amazon Simple Storage Service (S3) and Amazon Redshift in just a few clicks.

<https://aws.amazon.com/appflow/>

upvoted 5 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

<https://aws.amazon.com/appflow/>

"With Amazon AppFlow automate bi-directional data flows between SaaS applications and AWS services in just a few clicks."

If you want to pass the exam, choose B, regardless of your personal experience! Always use AWS managed services for "least operational overhead"

upvoted 6 times

[Removed] 1 year, 6 months ago

SaaS - AppFlow

upvoted 2 times

OmegaLambda7XL9 1 year, 6 months ago

Yea , I think this question is looking for Amazon Appflow.I also feel like it would be easier to set up Autoscaling for the already existing EC2 instances in the short term but then the fact that this software integrates with a lot of SAAS services means using Amazon Appflow will work reduce operational overhead in the long term

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #42

Topic 1


A company runs a highly available image-processing application on Amazon EC2 instances in a single VPC. The EC2 instances run inside several subnets across multiple Availability Zones. The EC2 instances do not communicate with each other. However, the EC2 instances download images from Amazon S3 and upload images to Amazon S3 through a single NAT gateway. The company is concerned about data transfer charges.

What is the MOST cost-effective way for the company to avoid Regional data transfer charges?

- A. Launch the NAT gateway in each Availability Zone.
- B. Replace the NAT gateway with a NAT instance.
- C. Deploy a gateway VPC endpoint for Amazon S3. **Most Voted**
- D. Provision an EC2 Dedicated Host to run the EC2 instances.

Correct Answer: C

Community vote distribution



Comments

SilentMilli **Highly Voted** 2 years, 5 months ago

Selected Answer: C

Deploying a gateway VPC endpoint for Amazon S3 is the most cost-effective way for the company to avoid Regional data transfer charges. A gateway VPC endpoint is a network gateway that allows communication between instances in a VPC and a service, such as Amazon S3, without requiring an Internet gateway or a NAT device. Data transfer between the VPC and the service through a gateway VPC endpoint is free of charge, while data transfer between the VPC and the Internet through an Internet gateway or NAT device is subject to data transfer charges. By using a gateway VPC endpoint, the company can reduce its data transfer costs by eliminating the need to transfer data through the NAT gateway to access Amazon S3. This option would provide the required connectivity to Amazon S3 and minimize data transfer charges.

upvoted 100 times

Bmarodi 2 years ago

Very good explanation!

upvoted 9 times

johne42 1 year, 9 months ago

<https://aws.amazon.com/blogs/storage/managing-amazon-s3-access-with-vpc-endpoints-and-s3-access-points/>

upvoted 5 times

OmegaLambda7XL9 1 year, 6 months ago

Precisely

upvoted 3 times

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: C

Option C (correct). Deploy a gateway VPC endpoint for Amazon S3.

A VPC endpoint for Amazon S3 allows you to access Amazon S3 resources within your VPC without using the Internet or a NAT gateway. This means that data transfer between your EC2 instances and S3 will not incur Regional data transfer charges.

Option A (wrong), launching a NAT gateway in each Availability Zone, would not avoid data transfer charges because the NAT gateway would still be used to access S3.

Option B (wrong), replacing the NAT gateway with a NAT instance, would also not avoid data transfer charges as it would still require using the Internet or a NAT gateway to access S3.

Option D (wrong), provisioning an EC2 Dedicated Host, would not affect data transfer charges as it only pertains to the physical host that the EC2 instances are running on and not the data transfer charges for accessing.

upvoted 6 times

Dharmarajan **Most Recent** 4 months, 2 weeks ago

Selected Answer: C

This one was obvious. with VPC gateway for S3, the traffic will not leave AWS and directly go to S3. Therefore will reduce the cost of outgoing data.

upvoted 1 times

satyaamm 5 months, 2 weeks ago

Selected Answer: C

Gateway VPC endpoint are most suitable for privately accessing S3 here.

upvoted 1 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: C

Using a gateway VPC endpoint for Amazon S3 is the most cost-effective solution because it enables instances within the VPC to communicate directly with Amazon S3 without incurring data transfer charges via a NAT gateway. Here's how it addresses the problem:

Avoids NAT Gateway Charges for S3 Access:

NAT gateways incur data processing and data transfer charges when instances access S3.

A gateway VPC endpoint eliminates the need to route traffic through the NAT gateway by enabling direct communication between EC2 instances and S3.

No Data Transfer Costs for Intra-Region S3 Access:

Data transferred between Amazon S3 and resources in the same AWS Region through a gateway VPC endpoint is free of charge.

Highly Available and Managed:

Gateway VPC endpoints are highly available and fully managed by AWS, requiring no additional operational overhead.

upvoted 1 times

atikla 5 months, 2 weeks ago

Selected Answer: C

Correct

upvoted 1 times

ChymKuBoy 8 months ago

Selected Answer: C

C for sure

upvoted 1 times

PaulGa 9 months ago

Selected Answer: C

Ans C - excellent explanation by SilentMilli

upvoted 1 times

monkey_aws 10 months ago

Selected Answer: C

VPC gateway endpoint is free to use, but only available for S3 and DynamoDB

upvoted 4 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

Gateway VPC allows direct access to S3 without going through public internet. This is the de-facto way to save cost for S3 to VPC traffic.

Correct answer is C

upvoted 3 times

[Removed] 1 year, 6 months ago

Avoid regional data transfer charge - VPC endpoint

upvoted 3 times

Ruffyit 1 year, 7 months ago

<https://aws.amazon.com/blogs/storage/managing-amazon-s3-access-with-vpc-endpoints-and-s3-access-points/>

upvoted 2 times

ACloud_Guru15 1 year, 8 months ago

Selected Answer: C

Gateway Endpoint best suits the requirement

upvoted 2 times

srinivasmn 1 year, 8 months ago

Answer is C: An S3 VPC endpoint provides a way for an S3 request to be routed through to the Amazon S3 service, without having to connect a subnet to an internet gateway. The S3 VPC endpoint is what's known as a gateway endpoint.

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

the EC2 instances are downloading and uploading images to S3, configuring a gateway VPC endpoint will allow them to access S3 without crossing Availability Zones or regions, eliminating regional data transfer charges

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

Gateway VPC endpoints provide reliable connectivity to Amazon S3 without requiring an internet gateway or a NAT device for your VPC.

upvoted 3 times

miki111 1 year, 11 months ago

Option C is the right answer.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #43

Topic 1


A company has an on-premises application that generates a large amount of time-sensitive data that is backed up to Amazon S3. The application has grown and there are user complaints about internet bandwidth limitations. A solutions architect needs to design a long-term solution that allows for both timely backups to Amazon S3 and with minimal impact on internet connectivity for internal users.

Which solution meets these requirements?

- A. Establish AWS VPN connections and proxy all traffic through a VPC gateway endpoint.
- B. Establish a new AWS Direct Connect connection and direct backup traffic through this new connection. **Most Voted**
- C. Order daily AWS Snowball devices. Load the data onto the Snowball devices and return the devices to AWS each day.
- D. Submit a support ticket through the AWS Management Console. Request the removal of S3 service limits from the account.

Correct Answer: B

Community vote distribution



Comments

Sinaneos **Highly Voted** 2 years, 8 months ago

Selected Answer: B

- A: VPN also goes through the internet and uses the bandwidth
- C: daily Snowball transfer is not really a long-term solution when it comes to cost and efficiency
- D: S3 limits don't change anything here

So the answer is B
upvoted 40 times

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: B

Option B (correct). Establish a new AWS Direct Connect connection and direct backup traffic through this new connection.

AWS Direct Connect is a network service that allows you to establish a dedicated network connection from your on-premises data center to AWS. This connection bypasses the public Internet and can provide more reliable, lower-latency communication between your on-premises application and Amazon S3. By directing backup traffic through the AWS Direct Connect connection, you can minimize the impact on your internet bandwidth and ensure timely backups to S3.

connection, you can minimize the impact on your internet bandwidth and ensure timely backups to S3.

upvoted 37 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option A (wrong), establishing AWS VPN connections and proxying all traffic through a VPC gateway endpoint, would not necessarily minimize the impact on internet bandwidth as it would still utilize the public Internet to access S3.

Option C (wrong), using AWS Snowball devices, would not address the issue of internet bandwidth limitations as the data would still need to be transferred over the Internet to and from the Snowball devices.

Option D (wrong), submitting a support ticket to request the removal of S3 service limits, would not address the issue of internet bandwidth limitations and would not ensure timely backups to S3.

upvoted 11 times

OmegaLambda7XL9 1 year, 6 months ago

Snowball isn't timely since it takes days after ordering to receive the Snowball devices and days to have it shipped and backed up

upvoted 3 times

Bofi 2 years, 3 months ago

Option C is wrong so is your reason. you do not need internet to load data into Snowball Devices. if you are using snow cone for example, u will connect it to your on-premises device directly for loading and Aws will load it in the cloud. However, it not effective to do that everyday , hence option B is the better choice.

upvoted 4 times

Buruguduystunstugudunstuy 2 years, 3 months ago

You're right Option B is the correct answer. I answered Option B as the correct answer above.

upvoted 1 times

Dharmarajan Most Recent 4 months, 2 weeks ago

Selected Answer: B

This one is clear from the given choices. There is no specification of "immediate" need, which means the company can afford to wait for that one month time period required to set up AWS Direct connect.

upvoted 1 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: B

Establishing an AWS Direct Connect connection provides a dedicated, high-bandwidth, low-latency link between the on-premises environment and AWS. This solution effectively addresses the problem by:

Reducing Internet Bandwidth Usage:

By routing backup traffic through the Direct Connect connection, internal users' internet bandwidth is preserved for other activities.

Ensuring Timely Backups:

Direct Connect provides a consistent and reliable connection with speeds ranging from 50 Mbps to 100 Gbps, ensuring that large volumes of time-sensitive data can be backed up efficiently.

Long-Term Scalability:

This solution supports long-term data transfer needs as the application's data volume grows

upvoted 1 times

Tjazz04 6 months ago

Selected Answer: B

Ans. B

For long term solution, using AWS Direct Connect for backup traffic will have its own bandwidth for backing-up data and will not share the bandwidth used by the users.

upvoted 1 times

zied007 9 months, 1 week ago

Selected Answer: B

Selected Answer: B

The answer is B
upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

Direct Connect is the only working solution.
A: VPN uses same bandwidth so doesn't solve anything
C: Snowball devices are physical devices requiring physical shipment so wrong solution
D: There are no S3 service limitations in the account related to this problem
upvoted 2 times

NicolasB 1 year, 5 months ago

Option B.
AWS Direct Connect link your on-premise instance with VPC, and all traffic will bypass your Internet Service Provider.
<https://docs.aws.amazon.com/directconnect/latest/UserGuide/Welcome.html>
upvoted 2 times

[Removed] 1 year, 6 months ago

Resolve Internet connection problem - Direct Connect
upvoted 2 times

Ruffyt 1 year, 7 months ago

AWS Direct Connect is a network service that allows you to establish a dedicated network connection from your on-premises data center to AWS. This connection bypasses the public Internet and can provide more reliable, lower-latency communication between your on-premises application and Amazon S3. By directing backup traffic through the AWS Direct Connect connection, you can minimize the impact on your internet bandwidth and ensure timely backups to S3.
upvoted 3 times

AWSSStudyBuddy 1 year, 7 months ago

Selected Answer: B

I picked option B, because AWS Direct Connect offers a dedicated, secure, high-performance connection that may circumvent bandwidth restrictions and minimize the impact on internet access, AWS Direct Connect is the ideal choice for backing up data to Amazon S3. Some solutions are not as good because they are not as scalable, reliable, or secure as VPN connections, Snowball devices, or reducing S3 service constraints.
upvoted 3 times

tom_cruise 1 year, 8 months ago

Selected Answer: B

Key: time sensitive. So snowball does not apply here.
upvoted 2 times

srinivasmn 1 year, 8 months ago

Right option is C,, In AWS Direct Connect, the network is not fluctuating and provides a consistent experience, while in AWS VPN the VPN is connected with shared and public networks, so the bandwidth and latency fluctuate. Hence direct connect is better choice than virtual connect.
upvoted 2 times

srinivasmn 1 year, 8 months ago

Typo correction to my my above comment. The right option is B.
upvoted 2 times

chandu7024 1 year, 8 months ago

Option B Correct. Reason is that, Direct connect will not use internet. But it will take good amount of time to establish the connectivity.
upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

AWS Direct Connect is a dedicated network connection between your on-premises network and AWS. This provides a private, high-bandwidth connection that is not subject to the same internet bandwidth limitations as traditional internet connections. This will allow for timely backups to Amazon S3 without impacting internet connectivity for internal users.

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: B

AWS Direct Connect cloud service is the shortest path to your AWS resources. While in transit, your network traffic remains on the AWS global network and never touches the public internet. This reduces the chance of hitting bottlenecks or unexpected increases in latency.

<https://aws.amazon.com/directconnect/#:~:text=The-,AWS%20Direct%20Connect,-cloud%20service%20is>

upvoted 2 times

miki111 1 year, 11 months ago

Option B is the right answer.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #44


Topic 1

A company has an Amazon S3 bucket that contains critical data. The company must protect the data from accidental deletion. Which combination of steps should a solutions architect take to meet these requirements? (Choose two.)

- A. Enable versioning on the S3 bucket. **Most Voted**
- B. Enable MFA Delete on the S3 bucket. **Most Voted**
- C. Create a bucket policy on the S3 bucket.
- D. Enable default encryption on the S3 bucket.
- E. Create a lifecycle policy for the objects in the S3 bucket.

Correct Answer: AB

Community vote distribution



Comments

Uhrien **Highly Voted** 2 years, 8 months ago

Selected Answer: AB

The correct solution is AB, as you can see here:

<https://aws.amazon.com/it/premiumsupport/knowledge-center/s3-audit-deleted-missing-objects/>

It states the following:

To prevent or mitigate future accidental deletions, consider the following features:

Enable versioning to keep historical versions of an object.

Enable Cross-Region Replication of objects.

Enable MFA delete to require multi-factor authentication (MFA) when deleting an object version.

upvoted 67 times

liams123 11 months, 1 week ago

but it could be C you could use 's3:deleteobject' permission without specific conditions or restricts only to authorized users. B does the same thing tho but B is mainly used to restrict unauthroized access not deletion. does anyone agree. I think it is A and C or A & B

upvoted 6 times

cookieMr Highly Voted 1 year, 11 months ago

Selected Answer: AB

Enabling versioning on S3 ensures multiple versions of object are stored in bucket. When object is updated or deleted, new version is created, preserving previous version.

Enabling MFA Delete adds additional layer of protection by requiring MFA device to be present when attempting to delete objects. This helps prevent accidental or unauthorized deletions by requiring extra level of authentication.

C. Creating a bucket policy on S3 is more focused on defining access control and permissions for bucket and its objects, rather than protecting against accidental deletion.

D. Enabling default encryption on S3 ensures that any new objects uploaded to bucket are automatically encrypted. While encryption is important for data security, it does not directly address accidental deletion.

E. Creating lifecycle policy for objects in S3 allows for automated management of objects based on predefined rules. While this can help with data retention and storage cost optimization, it does not directly protect against accidental deletion.

upvoted 13 times

liams123 11 months, 1 week ago

but it could be C you could use 's3:deleteobject' permission without specific conditions or restricts only to authorized users. B does the same thing tho but B is mainly used to restrict unauthroized access not deletion. does anyone agree. I think it is A and C or A & B

upvoted 2 times

Mrigranj12 Most Recent 4 months, 2 weeks ago

Selected Answer: AB

c: won't be able to delete anything

d: doesn't change anything

e: baseless

upvoted 1 times

satyaamm 5 months, 2 weeks ago

Selected Answer: AB

AB is the correct solution here as we need MFA delete for making the deletion process inevitable and also we need S3 bucket versioning for MFA delete.

upvoted 1 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: AB

To protect critical data in an S3 bucket from accidental deletion, these two features provide robust safeguards:

A. Enable versioning on the S3 bucket:

Versioning allows the bucket to maintain multiple versions of an object. Even if an object is accidentally deleted, the previous version can be restored, ensuring data recovery.

B. Enable MFA Delete on the S3 bucket:

MFA Delete adds an additional layer of security by requiring multi-factor authentication for deletion operations. This prevents accidental or unauthorized deletions.

upvoted 1 times

Chr1s_Mrg 8 months, 2 weeks ago

Selected Answer: AB

To protect critical data in an Amazon S3 bucket from accidental deletion, a solutions architect should take the following steps:

Enable versioning on the S3 bucket: This allows you to recover objects that are accidentally deleted or overwritten by keeping multiple versions of an object.

Enable MFA Delete on the S3 bucket: This adds an extra layer of security by requiring multi-factor authentication (MFA) for delete operations, which helps prevent accidental or unauthorized deletions

upvoted 2 times

PaulGa 9 months ago

Selected Answer: AB

Selected Answer: BU

Ans A,B - as per 'kwabsAA' 2 months ago

"To protect data from accidental deletion, the correct answers are B and D. Versioning does not prevent accidental deletion; it only allows for recovery after the fact. Multi-Factor Authentication (MFA) helps prevent accidental deletion by requiring an additional confirmation step before deletion, making it deliberate rather than accidental. Option D, which involves encryption, ensures that only individuals with the encryption keys can read or manipulate the data, thus preventing unauthorized access and manipulation, including deletion."

upvoted 2 times

zied007 9 months, 1 week ago

Selected Answer: AB

encryption will not prevent accidental deletions

upvoted 2 times

kwabsAA 11 months, 1 week ago

To protect data from accidental deletion, the correct answers are B and D. Versioning does not prevent accidental deletion; it only allows for recovery after the fact. Multi-Factor Authentication (MFA) helps prevent accidental deletion by requiring an additional confirmation step before deletion, making it deliberate rather than accidental. Option D, which involves encryption, ensures that only individuals with the encryption keys can read or manipulate the data, thus preventing unauthorized access and manipulation, including deletion.

upvoted 3 times

liams123 11 months, 1 week ago

you do realised B and D are only to stop unauthroized people from deleting it and if they did delete it would be on purpose to cause issues for a business. it says accidental id say A and C. idk the questions worded poorly dont trust half the answers on here

upvoted 2 times

liams123 11 months, 1 week ago

but it could be C you could use 's3:deleteobject" permission without specific conditions or restricts only to authorized users. B does the same thing tho but B is mainly used to restrict unauthroized access not deletion. does anyone agree. I think it is A and C or A & B. A allows multiple versions of objects to be stored in the bucket. Even if an object is deleted, its previous versions remain intact and accessible. idk the questions weird i could see how it could be all of them except D & E

upvoted 2 times

kwabsAA 11 months, 1 week ago

BD. For D, When you encrypt data, an unauthorized user (without the encryption key) cannot manipulate the data (ie. decryption, modifying, deletion).

upvoted 2 times

mmrakib 1 year, 3 months ago

Selected Answer: AB

AB will be the correct answer.

upvoted 3 times

sidharthwader 1 year, 3 months ago

This could be done if we enable MFA delete on the bucket but in order to enable this bucket versioning must be done. Hence A and B would be the answer.

upvoted 2 times

liams123 11 months, 1 week ago

but it could be C you could use 's3:deleteobject" permission without specific conditions or restricts only to authorized users. B does the same thing tho but B is mainly used to restrict unauthroized access not deletion. does anyone agree. I think it is A and C or A & B

upvoted 1 times

Conster 1 year, 4 months ago

I am getting so confused about what answers I should study. The answers don't match here or in ChatGPT. Can anyone who just took the exam, and passed, point me in the right direction? TIA!

upvoted 2 times

KRC96 1 year, 1 month ago

chatgpt will help you only if you gave correct prompt.

upvoted 1 times

liams123 11 months, 1 week ago

yea chatgpt said A and C it does make sense. Cause B is mainly used for unauthorized access not deletion. idk this website and certlibrary give some interesting answers makes it hard to know but some questions are difficult

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: AB

B: MFA to put an extra step to verify deletion and stop from accidental deletion

A: Versioning for recovery of objects that were deleted accidentally even with MFA

Remember, the solution is not required to STOP from deletion. It just wants to STOP ACCIDENTAL deletion.

CDE offer nothing related to accidental deletion

upvoted 3 times

rt_7777 1 year, 5 months ago

Not sure why Answer is BD. I am trying to rationalize it. What I guess could be to address keyword "critical data" where set default encryption is just enhance the security of stored data but does not prevent from deletion. This will be have 2 options A, B for that. B is make sense to ensure user know what to delete on second layer. For option A, it just help you to audit and recovered what was accidentally deleted but does not "prevent" accidentally delete.

upvoted 1 times

fb4afde 1 year, 6 months ago

Selected Answer: AB

Agree, s3 encryption does not prevent deletion

upvoted 2 times

jjcode 1 year, 6 months ago

Yeah so.. encryption is enabled by default on S3, soooooo why is the answer D.

Starting today, Amazon Simple Storage Service (Amazon S3) encrypts all new objects by default. Now, S3 automatically applies server-side encryption (SSE-S3) for each new object, unless you specify a different encryption option.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #45

Topic 1

A company has a data ingestion workflow that consists of the following:

- An Amazon Simple Notification Service (Amazon SNS) topic for notifications about new data deliveries
- An AWS Lambda function to process the data and record metadata

The company observes that the ingestion workflow fails occasionally because of network connectivity issues. When such a failure occurs, the Lambda function does not ingest the corresponding data unless the company manually reruns the job. Which combination of actions should a solutions architect take to ensure that the Lambda function ingests all data in the future? (Choose two.)

A. Deploy the Lambda function in multiple Availability Zones.

B. Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe it to the SNS topic. **Most Voted**


C. Increase the CPU and memory that are allocated to the Lambda function.

D. Increase provisioned throughput for the Lambda function.

E. Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue. **Most Voted**

Correct Answer: BE

Community vote distribution



Comments

Incognito013 **Highly Voted** 2 years, 8 months ago

A, C, D options are out, since Lambda is fully managed service which provides high availability and scalability by its own

Answers are B and E
upvoted 27 times

Oluseun 2 years, 3 months ago

There are times you do have to increase lambda memory for improved performance though. But not in this case.
upvoted 4 times

Sinaneos **Highly Voted** 2 years, 8 months ago

Selected Answer: BE

BE so that the lambda function reads the SQS queue and nothing gets lost
upvoted 12 times

AshishDhole Most Recent 5 months ago

Selected Answer: BE

best solution is to decouple the applications using SQS queue.
upvoted 1 times

Rcosmos 5 months ago

Selected Answer: BE

Explicação:
Criar uma fila do SQS e inscrevê-la no tópico do SNS (Opção B):

Quando o tópico do SNS envia notificações, elas são colocadas na fila do SQS.
Isso garante que as mensagens sejam armazenadas de forma durável na fila, mesmo que a função do Lambda não esteja disponível para processá-las imediatamente.
Modificar a função do Lambda para ler da fila do SQS (Opção E):

A função do Lambda é configurada para ser acionada por mensagens na fila do SQS.
Se ocorrer um problema de conectividade ou falha na função, as mensagens permanecerão na fila do SQS até que sejam processadas com sucesso.
Isso evita a perda de mensagens e elimina a necessidade de intervenção manual.
upvoted 1 times

satyaamm 5 months, 2 weeks ago

Selected Answer: BE

Enabling SQS with Lambda is the solution here as we need all data to be ingested in future.
upvoted 1 times

MGKYAING 5 months, 2 weeks ago

Selected Answer: BE

B. Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe it to the SNS topic:
By adding an SQS queue between the SNS topic and the Lambda function, you decouple the notification and processing layers. Messages are stored reliably in the SQS queue, and they are not lost even if the Lambda function fails temporarily due to network connectivity or other issues.

E. Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue:
Configuring the Lambda function to poll the SQS queue ensures that messages are processed asynchronously. If a processing failure occurs, the message remains in the queue and can be retried without manual intervention.
This ensures all data is eventually processed.
upvoted 3 times

friday_test 7 months ago

B,E : SQS is a pull mechanism broker, with at least one delivery, so by using SQS no messages get lost.
upvoted 1 times

Chr1s_Mrg 8 months, 2 weeks ago

Selected Answer: BE

Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe it to the SNS topic: This will decouple the data ingestion process and ensure that messages are not lost if the Lambda function fails to process them immediately

Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue: This ensures that the Lambda function can process the messages from the SQS queue, providing a reliable way to handle data ingestion even if there are network issues.
upvoted 1 times

ChymKuBoy 12 months ago

Selected Answer: BE

BE for sure
upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: BE

BE is correct as SQS ensures the messages are stored in a queue for processing.

A: No issue with Lambda availability so this solution is wrong

C: No issues with CPU or memory so no value added by this step also

D: This is not a provisioning issue so provisioning more Lambda won't solve the re-execution issues. The missed messages will still be lost

upvoted 4 times

OmegaLambda7XL9 1 year, 6 months ago

Since network timeout is the issue here, introduce SQS and read from it, that way when network goes down, data still remains in the queue and when connectivity is back, the lambda function can continue from the last data in the queue

upvoted 4 times

SaurabhTiwari1 9 months, 3 weeks ago

Nice explanation

upvoted 2 times

SaurabhTiwari1 9 months, 3 weeks ago

explanation*

upvoted 1 times

Ruffyit 1 year, 7 months ago

the correct combination of actions to ensure that the Lambda function ingests all data in the future is to create an SQS queue and subscribe it to the SNS topic (option B) and modify the Lambda function to read from the SQS queue (option E).

upvoted 1 times

tom_cruise 1 year, 8 months ago

Selected Answer: BE

Key: network connectivity issues

upvoted 2 times

awashenko 1 year, 8 months ago

Selected Answer: BE

This one told you the answer in the answer choices. Just add the word THEN between B and E and there ya go.

upvoted 1 times

Abdou1604 1 year, 10 months ago

B and E, the FAN out model, SQS will help to retrieve the work and delayed processing

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: BE

B) Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe it to the SNS topic.

E) Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue.

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: BE

BE is most logical answer.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions**Question #46**

Topic 1

A company has an application that provides marketing services to stores. The services are based on previous purchases by store customers. The stores upload transaction data to the company through SFTP, and the data is processed and analyzed to generate new marketing offers. Some of the files can exceed 200 GB in size.

Recently, the company discovered that some of the stores have uploaded files that contain personally identifiable information (PII) that should not have been included. The company wants administrators to be alerted if PII is shared again. The company also wants to automate remediation.


What should a solutions architect do to meet these requirements with the LEAST development effort?

A. Use an Amazon S3 bucket as a secure transfer point. Use Amazon Inspector to scan the objects in the bucket. If objects contain PII, trigger an S3 Lifecycle policy to remove the objects that contain PII.

B. Use an Amazon S3 bucket as a secure transfer point. Use Amazon Macie to scan the objects in the bucket. If objects contain PII, use Amazon Simple Notification Service (Amazon SNS) to trigger a notification to the administrators to remove the objects that contain PII. **Most Voted**

C. Implement custom scanning algorithms in an AWS Lambda function. Trigger the function when objects are loaded into the bucket. If objects contain PII, use Amazon Simple Notification Service (Amazon SNS) to trigger a notification to the administrators to remove the objects that contain PII.

D. Implement custom scanning algorithms in an AWS Lambda function. Trigger the function when objects are loaded into the bucket. If objects contain PII, use Amazon Simple Email Service (Amazon SES) to trigger a notification to the administrators and trigger an S3 Lifecycle policy to remove the objects that contain PII.

Correct Answer: B*Community vote distribution***Comments****Gatt** **Highly Voted** 2 years, 7 months ago

I have a problem with answer B. The question says: "automate remediation". B says that you inform the administrator and he removes the data manually, that's not automating remediation. Very weird, that would mean that D is correct - but it's so much harder to implement.

upvoted 48 times

ronaldchow 2 years, 5 months ago

remediation 2 years, 5 months ago

By "automate remediation", I thought it meant to use Amazon Macie to automate discovery on personally identifiable information.

<https://aws.amazon.com/macie/>

- Discover sensitive data across your S3 environment to increase visibility and automated remediation of data security risks.
upvoted 7 times

wamy1738 1 year, 1 month ago

The answer is B because it requires the "LEAST development effort".

The confusing part is that remediation is NOT automated because the solution alerts the administrators but still requires manual action. Its a bad question.

upvoted 7 times

3680113 1 year ago

B Correct, a bad and selfish question indeed. AWS doesn't care if they actually follow through on the standards they created, they are only interested in us answering the question right . MIND YOU ALL THE OPTIKNS REQUIRE MANUAL COMPLETION OF PROCESS.

upvoted 2 times

3680113 1 year ago

What kind of life cycle policy removes meat ?. Life cycle only removes objects that exceeds programmed time.

upvoted 2 times

Maxpayne009 2 years, 1 month ago

Macie has file size limit and clearly question mentions 200GB filesizes are possible. Lambda is the way to go ..

upvoted 7 times

pentium75 1 year, 5 months ago

You're confusing "files to retrieve samples from" with "files to analyze". Macie can analyze 20 GB files.

upvoted 5 times

Joxtat 2 years, 5 months ago

Pay attention to the entire question as in What should a solutions architect do to meet these requirements with the LEAST development effort? That is why Macie is used. Answer is B

upvoted 9 times

grzeev Highly Voted 2 years, 7 months ago

Selected Answer: B

Amazon Macie is a data security and data privacy service that uses machine learning (ML) and pattern matching to discover and protect your sensitive data

upvoted 24 times

grzeev 2 years, 7 months ago

Macie automatically detects a large and growing list of sensitive data types, including personally identifiable information (PII) such as names, addresses, and credit card numbers. It also gives you constant visibility of the data security and data privacy of your data stored in Amazon S3

upvoted 13 times

MundiChor Most Recent 2 months, 1 week ago

Selected Answer: B

Although I selected B on the account of "Least Developmental Effort", However I feel the question itself is wrong. In lieu of being a difficult and tricky question, it contradicts itself too many times.

1. Macie has a S3 file limitation of 20 GB (in case of ZIP archive), here the files are exceeding 200 GB with no indication of breaking the file or compression.

2. The question states that the customer wants to automate remediation; now the word remediation means "to solve" and automate remediation means that there should not be any manual effort. But none of the options provide any solution for that.

I would rate the trust of this question pretty low.

upvoted 1 times

MPG1970 2 months, 2 weeks ago

Selected Answer: C

Can't be B Macie has a 5Gb limit

C is better than D Lambda can be set up to remove data immediately and AWS state their #1 priority is security.

upvoted 1 times

LovaP 3 months, 1 week ago

Selected Answer: B

You can also leverage Macie integration with Amazon EventBridge and AWS Security Hub to monitor, process, and remediate findings by using other services, applications, and systems.

upvoted 1 times

Vandaman 3 months, 3 weeks ago

Selected Answer: D

The remediation must be automated

upvoted 1 times

Panknil 4 months, 1 week ago

Selected Answer: B

I think "The company wants administrators to be alerted if PII is shared again" is the key here.. so B is correct answer.

upvoted 1 times

kyd0nix 4 months, 1 week ago

Selected Answer: B

None are correct.

A does not alert the admins.

B does not automate solution.

C and D imply dev effort.

upvoted 1 times

AwsAbhiKumar 4 months, 2 weeks ago

Selected Answer: D

Macie has file limitation of 5 GB whereas Custom Lambda function allows you to handle files larger than 5 GB, overcoming Macie's limitation.

upvoted 1 times

Dharmarajan 4 months, 2 weeks ago

Selected Answer: B

B appears most appropriate out of the given option, however it does not address automation of remediation.

However in the view of remediation, the other options do not address it either, so B is most appropriate.

upvoted 1 times

darlar 4 months, 2 weeks ago

Selected Answer: D

The good answer is D, even if "sensitive data" = "macie".

Macie has limits for analyzing files, so 200GB won't pass :

Size of an individual file to analyze:

Adobe Portable Document Format (.pdf) file: 1,024 MB

Apache Avro object container (.avro) file: 8 GB

Apache Parquet (.parquet) file: 8 GB

Email message (.eml) file: 20 GB

GNU Zip compressed archive (.gz or .gzip) file: 8 GB

Microsoft Excel workbook (.xls or .xlsx) file: 512 MB

Microsoft Word document (.doc or .docx) file: 512 MB

Non-binary text file: 20 GB

TAR archive (.tar) file: 20 GB

ZIP compressed archive (.zip) file: 8 GB

Also, the answer B doesn't provide an auto-remediation, the admin still needs to remove the file manually

Also, the answer B doesn't provide an auto remediation, the admin still needs to remove the file manually.
Very tricky question, but I think the right fit is D.

upvoted 1 times

AshishDhole 5 months ago

Selected Answer: B

PII + S3 == Amazon Macie

upvoted 1 times

AwsAbhiKumar 4 months, 2 weeks ago

Well this is correct but not in this situation where some file can exceed 200 GB (Macie has limit of 5GB)

upvoted 1 times

Rcosmos 5 months ago

Selected Answer: B

Notificações com Amazon SNS:

Quando o Macie detecta PII, ele pode ser configurado para acionar uma notificação via SNS, alertando os administradores para tomar as ações necessárias.

Automação parcial:

Embora o Macie não remova automaticamente os objetos, ele permite que administradores sejam informados para realizar a correção manualmente, garantindo controle sobre os dados.

Mínimo esforço de desenvolvimento:

Essa abordagem utiliza serviços nativos da AWS sem necessidade de scripts ou funções personalizadas, reduzindo significativamente o tempo e o custo de implementação.

upvoted 1 times

satyaamm 5 months, 2 weeks ago

Selected Answer: B

Using Amazon Macie is most suitable for a S3 bucket and using SNS is also suitable as both of these services resolve the issues with least operational overhead.

upvoted 1 times

oluwafemiapara 5 months, 2 weeks ago

Selected Answer: B

Amazon macie is used to fish out pii

upvoted 1 times

Trevisan 6 months ago

Selected Answer: D

I would say the answer is D - Macie maximum file size is 20GB (If a file is larger than the applicable quota, Macie doesn't analyze any data in the file. according to AWS documentation:

<https://docs.aws.amazon.com/macie/latest/user/macie-quotas.html>

Also, B option doesn't meet the requirement of automate remediation.

upvoted 2 times

rudyydmirij 7 months, 3 weeks ago

B is closest, but Macie should trigger Lambda for remediation

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #47

Topic 1

A company needs guaranteed Amazon EC2 capacity in three specific Availability Zones in a specific AWS Region for an upcoming event that will last 1 week.

What should the company do to guarantee the EC2 capacity?

- A. Purchase Reserved Instances that specify the Region needed.
- B. Create an On-Demand Capacity Reservation that specifies the Region needed.
- C. Purchase Reserved Instances that specify the Region and three Availability Zones needed.
- D. Create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones needed. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Incognito013 **Highly Voted** 2 years, 8 months ago

Reserved instances are for long term so on-demand will be the right choice - Answer D
upvoted 35 times

simoneric88 1 year, 4 months ago

Confirmed! Reserved instances require a fixed one-year or three-year commitment. See <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html#capacity-reservations-differences>
upvoted 6 times

tch 2 months, 3 weeks ago

Amazon EC2 billing and purchasing options
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/instance-purchasing-options.html>
upvoted 1 times

tch 2 months, 3 weeks ago

Reserved Instances – Reduce your Amazon EC2 costs by making a commitment to a consistent instance configuration, including instance type and Region, for a term of 1 or 3 years.

Capacity Reservations – Reserve capacity for your EC2 instances in a specific Availability Zone

Capacity Reservations – Reserve capacity for your EC2 instances in a specific Availability Zone.
upvoted 1 times

Buruguduystunstugudunstuy Highly Voted 2 years, 5 months ago

Selected Answer: D

CORRECT

Option D. Create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones needed.

An On-Demand Capacity Reservation is a type of Amazon EC2 reservation that enables you to create and manage reserved capacity on Amazon EC2. With an On-Demand Capacity Reservation, you can specify the Region and Availability Zones where you want to reserve capacity, and the number of EC2 instances you want to reserve. This allows you to guarantee capacity in specific Availability Zones in a specific Region.

WRONG

Option A, purchasing Reserved Instances that specify the Region needed, would not guarantee capacity in specific Availability Zones.

Option B, creating an On-Demand Capacity Reservation that specifies the Region needed, would not guarantee capacity in specific Availability Zones.

Option C, purchasing Reserved Instances that specify the Region and three Availability Zones needed, would not guarantee capacity in specific Availability Zones as Reserved Instances do not provide capacity reservations.

upvoted 26 times

BlueVolcano1 2 years, 4 months ago

Another reason as to why Reserved Instances aren't the solution here is that you have to commit to either a 1 year or 3 year term, not 1 week.

upvoted 23 times

tch Most Recent 2 months, 3 weeks ago

Reserved Instances – Reduce your Amazon EC2 costs by making a commitment to a consistent instance configuration, including instance type and Region, for a term of 1 or 3 years.

Capacity Reservations – Reserve capacity for your EC2 instances in a specific Availability Zone.

do not confuse

upvoted 1 times

Rcosmos 5 months ago

Selected Answer: D

A Opção D (Reserva de capacidade sob demanda especificando a região e as zonas de disponibilidade) é a melhor solução, pois garante a capacidade necessária nas zonas de disponibilidade específicas para um evento de curta duração, atendendo ao requisito de forma eficiente e econômica.

upvoted 1 times

satyaamm 5 months, 2 weeks ago

Selected Answer: D

On-demand capacity instances in the required region and required availability zones is the most suitable here.

upvoted 1 times

friday_test 7 months ago

Reserved instances offer either a 1 year or 3 year term commitment. So on demand-instances would satisfy the need above. The answer is D, as the question also requires availability zones.

upvoted 1 times

Edawg173 1 year, 2 months ago

i always get these wrong

upvoted 2 times

TilTil 1 year, 2 months ago

Selected Answer: D

Picked C for this one and failed. Reserved Instances are reserved for 1-3 years so On-Demand Reservation makes more sense.

upvoted 3 times

JA2018 7 months, 1 week ago

Check out <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html#capacity-reservations-differences>

Terms for Capacity Reservations: No commitment required. Can be created and canceled as needed.

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html>

"On-Demand Capacity Reservations enable you to reserve compute capacity for your Amazon EC2 instances in a specific Availability Zone for any duration"

upvoted 2 times

awsgeek75 1 year, 5 months ago

Just to avoid any confusion, Reserved Instance also guarantee capacity reservation. However, the reason why we don't need Reserved Instances is because they are reserved for a duration of 1 to 3 years with a lock-in contract (no refunds!). The company is only interested in reserved capacity for a week so D is the best solution.

<https://aws.amazon.com/compare/the-difference-between-on-demand-instances-and-reserved-instances/>

upvoted 2 times

[Removed] 1 year, 6 months ago

Guarantee capacity on 3 AZ - on demand reservation, specify region & Availability Zone

upvoted 1 times

Ruffyit 1 year, 7 months ago

CORRECT

Option D. Create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones needed.

An On-Demand Capacity Reservation is a type of Amazon EC2 reservation that enables you to create and manage reserved capacity on Amazon EC2. With an On-Demand Capacity Reservation, you can specify the Region and Availability Zones where you want to reserve capacity, and the number of EC2 instances you want to reserve. This allows you to guarantee capacity in specific Availability Zones in a specific Region.

upvoted 2 times

awashenko 1 year, 8 months ago

Selected Answer: D

Reserved Instances have a commitment over a year so those are out. Option B only allows you to specify the Region and not the AZ.

Therefore, D is the only solution.

upvoted 2 times

Abdou1604 1 year, 10 months ago

its B , On-Demand Capacity Reservation allows you to reserve capacity for Amazon EC2 instances in a specific AWS Region, without specifying specific Availability Zones

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: D

D is the correct option to guarantee EC2 capacity in specific Availability Zones for a set timeframe.

On-Demand Capacity Reservations allow you to reserve EC2 capacity across specific Availability Zones for any duration. This guarantees you will have access to those resources.

upvoted 2 times

miki111 1 year, 11 months ago

Option D is the right answer.

upvoted 1 times

cookieMr 1 year, 11 months ago

Selected Answer: D

The most appropriate option to guarantee EC2 capacity in three specific Availability Zones in the desired AWS Region for the 1-week event is to create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones (option D).

A. Purchasing Reserved Instances that specify the Region needed does not guarantee capacity in specific Availability Zones.

B. Creating an On-Demand Capacity Reservation without specifying the Availability Zones would not guarantee capacity in the desired zones.

C. Purchasing Reserved Instances that specify the Region and three Availability Zones is not necessary for a short-term event and involves longer-term commitments.

upvoted 4 times

Abrar2022 2 years ago

Reserved instances is for long term

On-demand Capacity reservation enables you to choose specific AZ for any duration

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #48

Topic 1

A company's website uses an Amazon EC2 instance store for its catalog of items. The company wants to make sure that the catalog is highly available and that the catalog is stored in a durable location.

What should a solutions architect do to meet these requirements?

- A. Move the catalog to Amazon ElastiCache for Redis.
- B. Deploy a larger EC2 instance with a larger instance store.
- C. Move the catalog from the instance store to Amazon S3 Glacier Deep Archive.
- D. Move the catalog to an Amazon Elastic File System (Amazon EFS) file system. **Most Voted**

Correct Answer: D*Community vote distribution*

D (93%)

A (7%)

Comments**Six_Fingered_Jose** **Highly Voted** 2 years, 7 months ago**Selected Answer: D**

keyword is "durable" location
A and B is ephemeral storage
C takes forever so is not HA,
that leaves D
upvoted 52 times

Fakhrudin 1 year, 9 months ago

Yes, if you open EFS home page (<https://aws.amazon.com/efs/>), Amazon state, "Securely and reliably access your files with a fully managed file system designed for 99.99999999 percent (11 9s) durability and up to 99.99 percent (4 9s) of availability."
upvoted 13 times

rajendradba **Highly Voted** 2 years, 8 months ago**Selected Answer: D**

Elasticache is in Memory, EFS is for durability
upvoted 21 times

sirasdf 1 year, 3 months ago

That's wrong. EFS is highly available.

upvoted 4 times

MortisG Most Recent 1 month ago

Selected Answer: D

EFS because of durability and availability. Both initial requirements.

upvoted 1 times

Dharmarajan 4 months, 2 weeks ago

Selected Answer: D

Easy one. D. none of the other options are meeting the criteria of durability.

upvoted 1 times

satyaamm 5 months, 2 weeks ago

Selected Answer: D

EFS is most suitable for high availability and persistent storage.

upvoted 2 times

PaulGa 9 months ago

Selected Answer: D

Ans D - the clue is its the only option that offers "durable" (Not sure why the author suggests A??)

upvoted 2 times

jatric 11 months, 2 weeks ago

Selected Answer: D

Durability is the key differentiator

upvoted 2 times

uxrdbean 1 year ago

If they are in instance store now, would EFS maintain the performance necessary for them? I don't think the performance would be similar.

upvoted 1 times

sudohogan 1 year, 1 month ago

The ElastiCache option solves both storage and availability issues.

upvoted 1 times

HectorCosta 1 year, 1 month ago

Selected Answer: D

Cache is not Durable

upvoted 2 times

Edawg173 1 year, 2 months ago

D would be correct why would it give me A. :/

upvoted 2 times

TilTil 1 year, 2 months ago

Selected Answer: D

A and B are distractors.

D is durable and HA.

upvoted 2 times

modehqdah 1 year, 4 months ago

Selected Answer: D

durable location

upvoted 1 times

upvoted 1 times

arslantobe 1 year, 4 months ago

Option C, which suggests moving the catalog to Amazon S3 Glacier Deep Archive, is not a suitable choice for an active catalog that requires high availability and quick access. Glacier Deep Archive is designed for long-term archival and may not provide the low-latency access required for a catalog used in a website.

Therefore, option D is the most appropriate choice for achieving both high availability and durability for the catalog.

upvoted 4 times

vip2 1 year, 5 months ago

Selected Answer: A

SEE <https://aws.amazon.com/elasticache/faqs/#Redis>

upvoted 1 times

Hamso 1 year, 5 months ago

Redis can be made durable, supports failover and multi-AZ deployment, it's effective in catalog use cases. EFS is effective when a shared storage is needed

upvoted 5 times

rt_7777 1 year, 5 months ago

A and D, who win?

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #49

Topic 1

A company stores call transcript files on a monthly basis. Users access the files randomly within 1 year of the call, but users access the files infrequently after 1 year. The company wants to optimize its solution by giving users the ability to query and retrieve files that are less than 1-year-old as quickly as possible. A delay in retrieving older files is acceptable.

Which solution will meet these requirements MOST cost-effectively?

A. Store individual files with tags in Amazon S3 Glacier Instant Retrieval. Query the tags to retrieve the files from S3 Glacier Instant Retrieval.


B. Store individual files in Amazon S3 Intelligent-Tiering. Use S3 Lifecycle policies to move the files to S3 Glacier Flexible Retrieval after 1 year. Query and retrieve the files that are in Amazon S3 by using Amazon Athena. Query and retrieve the files that are in S3 Glacier by using S3 Glacier Select. **Most Voted**

C. Store individual files with tags in Amazon S3 Standard storage. Store search metadata for each archive in Amazon S3 Standard storage. Use S3 Lifecycle policies to move the files to S3 Glacier Instant Retrieval after 1 year. Query and retrieve the files by searching for metadata from Amazon S3.

D. Store individual files in Amazon S3 Standard storage. Use S3 Lifecycle policies to move the files to S3 Glacier Deep Archive after 1 year. Store search metadata in Amazon RDS. Query the files from Amazon RDS. Retrieve the files from S3 Glacier Deep Archive.

Correct Answer: B

Community vote distribution

**Comments**

masetromain **Highly Voted** 2 years, 8 months ago

Selected Answer: B

I think the answer is B:
Users access the files randomly

S3 Intelligent-Tiering is the ideal storage class for data with unknown, changing, or unpredictable access patterns, independent of object size or retention period. You can use S3 Intelligent-Tiering as the default storage class for virtually any workload, especially data lakes, data analytics, new applications, and user-generated content.

<https://aws.amazon.com/fr/s3/storage-classes/intelligent-tiering/>

invoted 45 times

upvoted 15 times

MutiverseAgent 1 year, 11 months ago

Agree, S3 Intelligent-Tiering meets all the requirements. The very important/crucial consideration here to satisfy that all files withing a year are instantly accessible is that the two options "Archive Access" and "Deep Archive Access" are not enabled in the "Archive rule actions" section present in the "Intelligent-Tiering Archive configurations" of the bucket. Those options are not enabled by default so this answer will work.

upvoted 3 times

sachin 2 years, 3 months ago

What about if the file you have not accessed 360 days and intelligent tier moved the file to Glacier and on 364 day you want to access the file instantly ?

I think C is right choice

upvoted 6 times

boringtangent 1 year, 2 months ago

bro u r forgetting cost effectiveness which is the requirement 1 yr in s3 standard will cost more than s3 intelligent tire.

upvoted 3 times

habibi03336 2 years, 3 months ago

It says "S3 Intelligent-Tiering is the ideal storage class for data with unknown, changing, or unpredictable access patterns". However, the statement says access pattern is predictable. It says there is frequent access about 1year.

upvoted 2 times

lofzee 1 year ago

Helps to read sometimes

upvoted 1 times

PaulEkwem 7 months, 3 weeks ago

Hey, don't be rude

upvoted 1 times

CarlosMarin 1 year, 1 month ago

It syas "... access the files INFREQUENTLY after 1 year"

upvoted 1 times

killbots 2 years, 2 months ago

it doesnt say predictable, it says files are accessed random. Random = Unpredictable. Answer is B

upvoted 12 times

ssoffline 2 years ago

Answer is C, why not intelligent Tiering

If the Intelligent-Tiering data transitions to Glacier after 180 days instead of 1 year, it would still be a cost-effective solution that meets the requirements.

With files stored in Amazon S3 Intelligent-Tiering, the data is automatically moved to the appropriate storage class based on its access patterns. In this case, if the data transitions to Glacier after 180 days, it means that files that are infrequently accessed beyond the initial 180 days will be stored in Glacier, which is a lower-cost storage option compared to S3 Standard.

upvoted 6 times

pentium75 1 year, 5 months ago

C involves moving "the files to S3 Glacier Instant Retrieval" which is not cost-effective since "a delay in retrieving older files is acceptable."

upvoted 4 times

RupeC 1 year, 11 months ago

With S3 Intelligent-Tiering, you can define rules that determine when objects should be moved from the frequent access tier

to the infrequent access tier, or vice versa, within S3 Standard storage classes.
upvoted 2 times

IngenieriaEGlobal 1 year, 8 months ago

The Answer is B. S3 Intelligent-Tiering stores objects in two access tiers: one tier that is optimized for frequent access and another lower-cost tier that is optimized for infrequent access. For a small monthly monitoring and automation fee per object, S3 Intelligent-Tiering monitors access patterns and moves objects that have not been accessed for 30 consecutive days to the infrequent access tier. There are no retrieval fees in S3 Intelligent-Tiering. If an object in the infrequent access tier is accessed later, it is automatically moved back to the frequent access tier. No additional tiering fees apply when objects are moved between access tiers within the S3 Intelligent-Tiering storage class. S3 Intelligent-Tiering is designed for 99.9% availability and 99.999999999% durability, and offers the same low latency and high throughput performance of S3 Standard
upvoted 7 times

Visinho 1 year, 4 months ago

Are you not going to pay for Athena usage?
upvoted 1 times

Lilibell **Highly Voted** 2 years, 8 months ago

The answer is B
upvoted 12 times

c12ab95 **Most Recent** 1 month ago

Selected Answer: B

Why Option B Works Best:
First Year (Active Access):

S3 Intelligent-Tiering automatically shifts objects between Frequent and Infrequent Access tiers, reducing costs by ~40% if access declines.

Athena queries data directly in S3, avoiding retrieval fees for active tiers.

After 1 Year (Infrequent Access):

Lifecycle policies transition data to Glacier Flexible Retrieval, which supports Glacier Select for SQL-based partial retrievals (avoids restoring entire files).

Cost savings: Flexible Retrieval is cheaper than S3 Standard-IA and allows expedited (1–5 min) or bulk (5–12 hr) retrievals.

Query Flexibility:

Recent files (<1 year) are queried via Athena with millisecond latency.

Older files (>1 year) use Glacier Select to retrieve subsets of data, minimizing retrieval costs and time.
upvoted 1 times

sammo08 3 months, 3 weeks ago

Selected Answer: B

Option A Incorrect – S3 Glacier Instant Retrieval is more expensive than necessary for frequently accessed files, and tag-based querying is inefficient.

Option C Incorrect – Using S3 Glacier Instant Retrieval after 1 year is more expensive than needed, and storing metadata in S3 is less efficient than Athena.

Option D Incorrect – S3 Glacier Deep Archive has long retrieval times, and maintaining metadata in RDS adds extra costs.

Option B Correct – S3 Intelligent-Tiering optimizes costs, and S3 Glacier Flexible Retrieval balances cost with acceptable retrieval delays, while Athena and S3 Glacier Select enable efficient queries.
upvoted 1 times

dariar 4 months, 2 weeks ago

Selected Answer: B

The B and C are both good options, but since Athena is free for use with S3, it definitely the B-option.

Athena queries data directly from Amazon S3. There are no additional storage charges for querying your data with Athena.

upvoted 1 times

ricktechie66 5 months, 1 week ago

Selected Answer: C

This solution best meets the requirements while being cost-effective because:

1. **Initial storage in S3 Standard:** Provides the fastest access for files less than 1 year old, meeting the requirement for quick retrieval of newer files
2. **Metadata storage in S3 Standard:** Enables efficient querying and searching of files without additional database costs
3. **Lifecycle policy to S3 Glacier Instant Retrieval:** Provides cost optimization for older files while still maintaining relatively quick retrieval when needed.

upvoted 1 times

Rohan_Butala 5 months, 4 weeks ago

Selected Answer: B

Option B strikes the right balance between cost-effectiveness and performance, ensuring that frequently accessed data is easily retrievable and older data is archived at lower cost with fast retrieval options for both new and older data.

upvoted 1 times

_adoo 7 months, 1 week ago

If you are like me and thought Athena would make this cost-inefficient:

Athena is free for s3 querying.
It has a cost per TB for SQL Queries.
And a different cost for Apache apps.

<https://aws.amazon.com/athena/pricing/>

upvoted 3 times

JonesNick 7 months, 1 week ago

Keyword : MOST cost-effectively
Athena is great tool for analyzing data in S3. But it comes with the cost.
So answer is C.

upvoted 1 times

jatric 11 months, 2 weeks ago

Selected Answer: B

cost effective to retrieve the file of 1 year or less. Standard S3 is more cost effective than intelligent tiering.

upvoted 2 times

344bba0 11 months, 3 weeks ago

Selected Answer: C

The answer is C.

- * Intelligent tiering may not guarantee frequent, but random access and fast searches within a year.
- * Athena is a great analysis solution, but it is an unnecessary cost for search purposes only.

upvoted 4 times

mknarula 11 months, 3 weeks ago

Answer is C. The only difference between choice B and C is Glacier Storage class. The question states clearly that "giving users the ability to query and retrieve files that are less than 1-year-old as quickly as possible". This is possible via instant retrieval and not flexible retrieval.

upvoted 2 times

ManikRoy 1 year, 1 month ago

Selected Answer: B

option B is the only one that mentions use of Amazon Athena and Glacier select for querying, So I'll go with it though I would have preferred using S3 standard storage in place of S3 intelligent tiering for the first year.

upvoted 1 times

hro 1 year, 2 months ago

Answer is C - unfortunately. S3 Glacier Instant Retrieval - MOST Cost effective deep storage for retrieving files as quickly as

...then is a... immediately, so Glacier instant retrieval... then... ensure deep storage for retrieving files as quickly as possible.

upvoted 1 times

Liorage 1 year, 2 months ago

Selected Answer: B

B is good answer: Athena is good option to query data in S3. And before 1 year data are randomly use, for this, intelligent tiering is good option.

upvoted 1 times

chickenmf 1 year, 3 months ago

Selected Answer: C

The company wants to optimize its solution by giving users the ability to query and retrieve files that are less than 1-year-old as quickly as possible" -- What if S3 Intelligent-Tiering transitioned the data that's under 1 year old into a storage class that takes a long time to access?

upvoted 1 times

dsshahu01 1 year, 3 months ago

Selected Answer: C

The answer is C

S3 standard for first part (intelligent tiering is much better and cost-effective)
glacier instant retrieval because of the statement after an year needs to retrieved as soon as possible)

Also why ruling out B is because of Athena - it becomes expensive if data is retrieved using it after scanning all the data in glacier per request.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #50

Topic 1

A company has a production workload that runs on 1,000 Amazon EC2 Linux instances. The workload is powered by third-party software. The company needs to patch the third-party software on all EC2 instances as quickly as possible to remediate a critical security vulnerability.

What should a solutions architect do to meet these requirements?

- A. Create an AWS Lambda function to apply the patch to all EC2 instances.
- B. Configure AWS Systems Manager Patch Manager to apply the patch to all EC2 instances.
- C. Schedule an AWS Systems Manager maintenance window to apply the patch to all EC2 instances.
- D. Use AWS Systems Manager Run Command to run a custom command that applies the patch to all EC2 instances.

Most Voted**Correct Answer: D***Community vote distribution***Comments****tinyfoot** **Highly Voted** 2 years, 7 months ago

The primary focus of Patch Manager, a capability of AWS Systems Manager, is on installing operating systems security-related updates on managed nodes. By default, Patch Manager doesn't install all available patches, but rather a smaller set of patches focused on security. (Ref <https://docs.aws.amazon.com/systems-manager/latest/userguide/patch-manager-how-it-works-selection.html>)

Run Command allows you to automate common administrative tasks and perform one-time configuration changes at scale. (Ref <https://docs.aws.amazon.com/systems-manager/latest/userguide/execute-remote-commands.html>)

Seems like patch manager is meant for OS level patches and not 3rd party applications. And this falls under run command wheelhouse to carry out one-time configuration changes (update of 3rd part application) at scale.

upvoted 71 times

Fakhrudin 1 year, 9 months ago

3rd party applications are also supported by Patch Manager (<https://docs.aws.amazon.com/systems-manager/latest/userguide/patch-manager.html>).

You can use Patch Manager to apply patches for both operating systems and applications. (On Windows Server, application support is limited to updates for applications released by Microsoft.) You can use Patch Manager to install Service Packs on Windows nodes and perform minor version upgrades on Linux nodes. You can patch fleets of Amazon Elastic Compute Cloud

windows nodes and perform minor version upgrades on Linux nodes. You can patch fleets of Amazon Elastic Compute Cloud (Amazon EC2) instances, edge devices, on-premises servers, and virtual machines (VMs) by operating system type. This includes supported versions of several operating systems, as listed in Patch Manager prerequisites.

upvoted 8 times

Shasha1 Highly Voted 2 years, 6 months ago

D
AWS Systems Manager Run Command allows the company to run commands or scripts on multiple EC2 instances. By using Run Command, the company can quickly and easily apply the patch to all 1,000 EC2 instances to remediate the security vulnerability.

Creating an AWS Lambda function to apply the patch to all EC2 instances would not be a suitable solution, as Lambda functions are not designed to run on EC2 instances. Configuring AWS Systems Manager Patch Manager to apply the patch to all EC2 instances would not be a suitable solution, as Patch Manager is not designed to apply third-party software patches. Scheduling an AWS Systems Manager maintenance window to apply the patch to all EC2 instances would not be a suitable solution, as maintenance windows are not designed to apply patches to third-party software

upvoted 22 times

sammo08 Most Recent 3 months, 3 weeks ago

Selected Answer: D

Option A Incorrect – AWS Lambda is not suitable for managing and patching third-party software on EC2 instances at scale.

Option B Incorrect – AWS Systems Manager Patch Manager is designed for OS patching, not third-party software patching

Option C Incorrect – A maintenance window schedules updates but does not immediately apply the patch, delaying remediation

Option D Correct – AWS Systems Manager Run Command allows immediate execution of a custom patching script across all EC2 instances, ensuring quick remediation.

upvoted 3 times

Dharmarajan 4 months, 2 weeks ago

Selected Answer: B

I would vote for B as the option for Run command does not have exception handling capabilities like patch failures. OR the command has to be scripted to handle such cases, which is additional operational overhead - especially for cases where the exception corner cases are not handled in the script.

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: D

Patch Manager is optimized for OS-level security patches and may not support the specific third-party software patch required in this scenario.

Even if the third-party software is supported, Patch Manager may not provide the level of control or speed needed for a critical security update.

upvoted 1 times

Rcosmos 5 months ago

Selected Answer: D

A Opção D (AWS Systems Manager Run Command) é a solução mais rápida e eficaz para aplicar o patch a todas as 1.000 instâncias EC2 em resposta a uma vulnerabilidade crítica de segurança, garantindo uma correção imediata e controlada.

upvoted 1 times

Rohan_Butala 5 months, 4 weeks ago

Selected Answer: B

AWS Systems Manager Patch Manager is a fully managed service designed to automate the process of patching operating systems and applications on EC2 instances. It can automatically apply patches to EC2 instances at scale, ensuring that critical vulnerabilities are remediated quickly. Patch Manager allows you to define patch baselines, schedule patching, and apply patches across multiple instances without manual intervention.

upvoted 2 times

Rohan_Butala 5 months, 4 weeks ago

Selected Answer: B

Option B has right balance between cost-effectiveness and performance, ensuring that frequently accessed data is easily retrievable and older data is archived at lower cost with fast retrieval options for both new and older data.

upvoted 2 times

Homine 7 months, 4 weeks ago

Selected Answer: B

Path Manager can support patch third party applications

upvoted 2 times

PaulGa 9 months ago

Selected Answer: B

Ans B (should be). Clue: "The company needs to *patch* the third-party software..." - if the Patch Manager is not capable of this then perhaps it should be, rather than delegating functionality to another service... The suggested answer D, implies Patch Manager can't do the job...

upvoted 3 times

jhoiti 11 months, 3 weeks ago

Selected Answer: B

In your case, since the company needs to fix a critical security vulnerability as quickly as possible, Patch Manager (option B) would be the most recommended choice. It provides a quick and efficient way to apply the patch to all affected instances without the need to create custom scripts.

However, if your company has specific requirements that cannot be met by Patch Manager, Run Command (option D) may be a viable alternative.

upvoted 2 times

lofzee 1 year ago

I originally thought B but after a bit of reading I've changed my mind to D purely because patch manager will not be aware of this random third party application.

Run command allows you to install applications, run powershell, bash etc commands at scale so the most sensible answer would be run command.

upvoted 6 times

EMPERBACH 1 year, 1 month ago

Selected Answer: B

Install software -> Patch Manager

Run command/processing workload -> Run Command

upvoted 3 times

ManikRoy 1 year, 1 month ago

Selected Answer: D

I think patch manager would need an agent to be installed and also Patch Manager doesn't derive severity levels from third-party sources.

upvoted 3 times

RafikTAAMMA 1 year, 2 months ago

Selected Answer: D

AWS Systems Manager Patch Manager primarily focuses on operating system patches and does not directly support third-party software patching on Linux instances

upvoted 4 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

Critical means immediate. Just run the patch command with AWS SM run command to get it done. D is best choice.

A: Too convoluted

B: Can work but have to setup a lot of things to get this done. would be a good choice if D wasn't an option

C: It's a critical patch so not time for maintenance window

upvoted 5 times

rt_7777 1 year, 5 months ago

By practice, isn't schedule planned downtime is common sense before patching done?
upvoted 1 times

JA2018 7 months, 1 week ago

it will depend on the severity (blast radius) of the vulnerability declared.
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #51

Topic 1

A company is developing an application that provides order shipping statistics for retrieval by a REST API. The company wants to extract the shipping statistics, organize the data into an easy-to-read HTML format, and send the report to several email addresses at the same time every morning.

Which combination of steps should a solutions architect take to meet these requirements? (Choose two.)

A. Configure the application to send the data to Amazon Kinesis Data Firehose.

B. Use Amazon Simple Email Service (Amazon SES) to format the data and to send the report by email. **Most Voted**


C. Create an Amazon EventBridge (Amazon CloudWatch Events) scheduled event that invokes an AWS Glue job to query the application's API for the data.

D. Create an Amazon EventBridge (Amazon CloudWatch Events) scheduled event that invokes an AWS Lambda function to query the application's API for the data. **Most Voted**

E. Store the application data in Amazon S3. Create an Amazon Simple Notification Service (Amazon SNS) topic as an S3 event destination to send the report by email.

Correct Answer: BD

Community vote distribution



Comments

whosawsome **Highly Voted** 2 years, 8 months ago

Selected Answer: BD

You can use SES to format the report in HTML.
<https://docs.aws.amazon.com/ses/latest/dg/send-email-formatted.html>
 upvoted 33 times

apchandana 2 years, 2 months ago

this document is talking about the SES API. not ses. SES does not format data. just sending emails.
<https://aws.amazon.com/ses/>
 upvoted 7 times

Clouddon 1 year, 10 months ago

When you send an email with Amazon SES, the email information you need to provide depends on how you call Amazon SES. You can provide a minimal amount of information and have Amazon SES take care of all of the formatting for you. Or, if you want to do something more advanced like send an attachment, you can provide the raw message yourself.
<https://docs.aws.amazon.com/ses/latest/dg/send-email-concepts-email-format.html>

upvoted 4 times

backbencher2022 Highly Voted 2 years, 7 months ago

Selected Answer: BD

B&D are the only 2 correct options. If you are choosing option E then you missed the daily morning schedule requirement mentioned in the question which cant be achieved with S3 events for SNS. Event Bridge can used to configure scheduled events (every morning in this case). Option B fulfills the email in HTML format requirement (by SES) and D fulfills every morning schedule event requirement (by EventBridge)

upvoted 26 times

PaulGa 9 months ago

Option B doesn't specify an API interface...

upvoted 1 times

RupeC 1 year, 11 months ago

I don't believe you are correct when you say that E cannot meet the scheduling requirement. If the glue action is scheduled and outputs to S3, then as the S3 event destination is SNS, in effect you have a way of getting SNS to have a scheduled release.

upvoted 2 times

pentium75 1 year, 5 months ago

But E does not include a glue action. We need either C or D for the scheduling, plus B or E for the email sending.

upvoted 2 times

slimen 1 year, 7 months ago

the daily schedule can be achieve with event bridge

- schedule and event bridge to trigger daily
- the event briodge will trigger a lambda function that will collect data and save it in s3
- once data in s3 the event noitification will trigger SNS to send emails

upvoted 2 times

RamanadhRavinuthala Most Recent 10 months, 3 weeks ago

Selected Answer: BD

D Explanation: EventBridge can be used to schedule regular invocations of a Lambda function that retrieves the required data from the application's API. This step sets up the process to collect the data at the specified time every morning.

B Explanation: Amazon SES can format the data into an easy-to-read HTML report and send the email to multiple recipients efficiently.

upvoted 4 times

JA2018 7 months, 1 week ago

<https://docs.aws.amazon.com/ses/latest/dg/send-email-concepts-email-format.html>

upvoted 1 times

hieunt.hus 11 months ago

Selected Answer: BC

SNS & Glue

upvoted 5 times

jatric 11 months, 2 weeks ago

BC - Question says extract the data - which is the job of AWS Glue and then invoke SES endpoint to send email. SES support html format.

But BD also seems possible.

upvoted 5 times

lofzee 1 year ago

Defo B and D

Look at the aws notes on Sns

Notes

You can't customize the body of the email message. The email delivery feature is intended to provide internal system alerts, not marketing messages.

It can't send anything but notifications. SES can send normal emails therefore, BD is the answer.

upvoted 3 times

jhoiti 1 year, 2 months ago

Selected Answer: BE

Option B suggests using Amazon SES, which allows you to format the data and email the report to multiple recipients in an efficient and scalable way.

Option E proposes storing application data in Amazon S3, which is scalable and durable storage. By configuring an Amazon SNS topic as an S3 event destination, you can automatically trigger the report to be sent via email whenever new data is added to S3.

upvoted 1 times

Monster07 1 year, 2 months ago

Selected Answer: BD

Chat GPT says :

With Amazon SES, you can send rich, formatted email content, including text, HTML, attachments, and embedded images, suitable for email communication.

Amazon SNS is primarily used for sending plain-text or JSON-formatted messages, suitable for notifications and alerts across different channels.

This can suggest that we need to use SES if we want to use HTML content.

upvoted 1 times

wyejay 1 year, 5 months ago

Answer: B and D

Other options

A. Amazon Kinesis Data Firehose: This service is typically used for real-time streaming data processing rather than for scheduled tasks like generating a morning report.

C. Amazon EventBridge to invoke an AWS Glue job: AWS Glue is a data integration service that's more focused on ETL (extract, transform, load) operations, often involving large datasets and complex transformations, which might be more than needed for this scenario.

E. Amazon S3 with SNS topic: Storing data in S3 and using SNS for notification is viable, but this doesn't directly address the need to format the data into HTML and send it as an email report. SNS is better suited for sending notifications rather than formatted reports.

upvoted 10 times

tonybuivannghia 8 months, 4 weeks ago

Some mistakes that Kinesis Data Firehose is used for near real-time streaming data processing. But totally I agree with you that B & D are correct.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: BD

Very detailed question so let's break it down:

"send the report to several email addresses at the same time every morning" this locks B as nothing else can do it.

A: Firehose to collect data from API will work but it cannot generate a report

C: Glue is ETL, it cannot extract data from an API

E: Store data in S3. No idea what this will help with

The API provides order shipping data so you can query it. Lambda can be used to query the API easily so D is good choice that works with B.

BD is correct combination

upvoted 3 times

upvoted 5 times

pentium75 1 year, 5 months ago

Selected Answer: BD

"At the same time every morning" requires scheduling, which is only mentioned in C and D. AWS Glue has no native functionality to query REST APIs, thus we need a Lambda function -> D.

For email we need SES or SNS, but as we want "an easy-to-read HTML format", SNS is out. SNS can send notifications, not formatted emails. Thus B.

upvoted 5 times

[Removed] 1 year, 6 months ago

Key: Send email every morning same time - 1. Simple email 2. AWS Event Bridge with lambda

upvoted 2 times

wearrexdzw3123 1 year, 7 months ago

Selected Answer: B

I think there is a problem with the answer. It should be that ses sends the email processed by lambda.

upvoted 1 times

tom_cruise 1 year, 7 months ago

Selected Answer: BD

Key: retrieval by a REST API, that's why use lambda

upvoted 2 times

tom_cruise 1 year, 8 months ago

Selected Answer: DE

Both SES and SNS can format html, but there is a disconnection between B and D. Where do you store the data between the steps?

upvoted 4 times

pentium75 1 year, 5 months ago

Why would I need to "store the data"? Wouldn't the Lambda function just call the SES API?

upvoted 2 times

David_Ang 1 year, 8 months ago

Selected Answer: BD

the reason why "B" is more correct than "E" is because is more simple and you don't have to store data is not what they want, also SES is a service that is meant for sending the data through email, and is exactly what the company wants. is not the first time the admin is wrong with the answer

upvoted 3 times

hieulam 1 year, 8 months ago

Selected Answer: DE

E should be correct:

<https://saturncloud.io/blog/how-to-send-html-mails-using-amazon-sns/>

upvoted 1 times

h_sahu 1 year, 8 months ago

I believe BD are the answers. E can't be used, because, in E can't help with email formatting. E won't be the best choice even for scheduling.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #52

Topic 1

A company wants to migrate its on-premises application to AWS. The application produces output files that vary in size from tens of gigabytes to hundreds of terabytes. The application data must be stored in a standard file system structure. The company wants a solution that scales automatically, is highly available, and requires minimum operational overhead. Which solution will meet these requirements?

- A. Migrate the application to run as containers on Amazon Elastic Container Service (Amazon ECS). Use Amazon S3 for storage.
- B. Migrate the application to run as containers on Amazon Elastic Kubernetes Service (Amazon EKS). Use Amazon Elastic Block Store (Amazon EBS) for storage.
- C. Migrate the application to Amazon EC2 instances in a Multi-AZ Auto Scaling group. Use Amazon Elastic File System (Amazon EFS) for storage. **Most Voted**
- D. Migrate the application to Amazon EC2 instances in a Multi-AZ Auto Scaling group. Use Amazon Elastic Block Store (Amazon EBS) for storage.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**ArielSchivo** **Highly Voted** 2 years, 7 months ago**Selected Answer: C**

EFS is a standard file system, it scales automatically and is highly available.
upvoted 33 times

masetromain **Highly Voted** 2 years, 8 months ago

I have absolutely no idea...

Output files that vary in size from tens of gigabytes to hundreds of terabytes

Size limit for a single object:

S3 5 To TiB

<https://aws.amazon.com/fr/blogs/aws/amazon-s3-object-size-limit/>

EBS 64 TiB

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/volume_constraints.html

EFS 47.9 TiB

<https://docs.aws.amazon.com/efs/latest/ug/limits.html>

upvoted 11 times

RBSK 2 years, 6 months ago

None meets 100s of TB / file. Bit confusing / misleading

upvoted 6 times

Help2023 2 years, 3 months ago

The answer to that is

Limit size for a single object:

S3, 5TiB is per object but you can have more than one object in a bucket, meaning infinity

<https://aws.amazon.com/fr/blogs/aws/amazon-s3-object-size-limit/>

EBS 64 Tib is per block of storage

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/volume_constraints.html

EFS 47.9 TiB per file and in the questions its says Files the 's'

<https://docs.aws.amazon.com/efs/latest/ug/limits.html>

upvoted 2 times

JayBee65 2 years, 6 months ago

S3 and EBS are block storage but you are looking to store files, so EFS is the correct option.

upvoted 4 times

Ello2023 2 years, 5 months ago

S3 is object storage.

upvoted 14 times

OmegaLambda7XL9 1 year, 6 months ago

A lil correction,S3 is Object storage not Block Storage

upvoted 4 times

SirDNS Most Recent 2 months, 3 weeks ago

Selected Answer: C

C: The only option with file system

upvoted 1 times

AwsAbhiKumar 4 months, 2 weeks ago

Selected Answer: C

HOW EFS is suitable here as it has a maximum file size limit of 52 TiB, so it cannot store a single file that exceeds this size. But in question it has mentioned that output files may vary in size from tens of gigabytes to hundreds of terabytes. Isn't it contradictory?

upvoted 1 times

Dharmarajan 4 months, 2 weeks ago

Selected Answer: C

The question is not clear. None of the options suit the 100s of TB requirements.

EFS max file size is clearly specified as 49.5TB.

But lets assume for a minute that the app is capable of multipart files, then EFS would be a great fit. (requirement: Must be a file system)

upvoted 1 times

satyaamm 5 months, 1 week ago

Selected Answer: C

EFS is the only file system structure here so it it most suitable.

upvoted 1 times

trinh_le 6 months, 4 weeks ago

Selected Answer: C

Keyword

Must store in standard file system -> ignore A

Minimum operational -> ignore B

Scalable + Terabytes-> pick C

upvoted 1 times

PaulGa 9 months ago

Selected Answer: C

Ans C - S3 unlimited storage in both scope and size of individual objects

upvoted 1 times

bishtr3 11 months ago

C : EFS as It is built to scale on demand to petabytes without disrupting applications, growing and shrinking automatically as you add and remove files.

Multiple compute instances, including Amazon EC2, Amazon ECS, and AWS Lambda, can access an Amazon EFS file system at the same time, providing a common data source for workloads.

upvoted 2 times

HectorCosta 1 year, 1 month ago

Selected Answer: C

Key words: Standard File System and Scales Automatically.

S3 is object Store, hence it fails with the "Standard File System" requirement, so we can discard A.

EBS does not scale automatically, failing with the "Scales Automatically" requirement, so we can discard B and D

upvoted 3 times

sidharthwader 1 year, 3 months ago

C is the only option which supports standard file system when we talk about high availability. EBS scope is within a availability zone but EFS has scope of a region.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

Standard file system that is highly available: EFS

Autoscaling highly available system: EC2 or ECS or EKS can work

A: Not suitable due to S3 which is BLOB not file system

B: EKS is ok but EBS is not HA

D: EBS is not HA

So by elimination, C is best option.

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: C

"File system structure" = EFS, which also meets all the other requirements.

upvoted 4 times

Mikado211 1 year, 6 months ago

Selected Answer: C

Technically the A could work, ECS is often recommended by AWS in case of minimum operational overhead, and S3 is durable and highly scalable BUT it is not a "traditional" file system structure. In an S3 bucket, there is no real file structure, only files and prefixes that simulate a structure.

B is wrong because of EKS which require more management

EFS is recommended for minimum operational overhead instead of EBS.

So C (EC2 + EFS) is recommended here over D (EC2 + EBS).

upvoted 4 times

manish123 1 year, 6 months ago