

wantu 1 year, 6 months ago

Selected Answer: C

Palabras clave: autoescalado y ficheros

upvoted 1 times

leosmal 1 year, 6 months ago

The key is Multi-AZ ,EBS does not support it.

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

Standard file system structure, scales automatically, requires minimum operational overhead = Amazon Elastic File System (Amazon EFS)

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #53

Topic 1

A company needs to store its accounting records in Amazon S3. The records must be immediately accessible for 1 year and then must be archived for an additional 9 years. No one at the company, including administrative users and root users, can be able to delete the records during the entire 10-year period. The records must be stored with maximum resiliency. Which solution will meet these requirements?

- A. Store the records in S3 Glacier for the entire 10-year period. Use an access control policy to deny deletion of the records for a period of 10 years.
- B. Store the records by using S3 Intelligent-Tiering. Use an IAM policy to deny deletion of the records. After 10 years, change the IAM policy to allow deletion.
- C. Use an S3 Lifecycle policy to transition the records from S3 Standard to S3 Glacier Deep Archive after 1 year. Use S3 Object Lock in compliance mode for a period of 10 years. **Most Voted**
- D. Use an S3 Lifecycle policy to transition the records from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 1 year. Use S3 Object Lock in governance mode for a period of 10 years.

Correct Answer: C

Community vote distribution

C (100%)

Comments

awsgeek75 **Highly Voted** 1 year, 5 months ago

Selected Answer: C

Only CD provides Object Lock options which is required for stopping admin/root users from deleting.
D is governance mode which is like government, pay enough money and you can do anything. This is not what we want so compliance is the option.
C is right choice.

For future, remember
S3 Lock Governance = corrupt government official
S3 Lock Compliance = honest solution architect!
upvoted 22 times

Guru4Cloud **Highly Voted** 1 year, 10 months ago

Selected Answer: C

The key reasons are:

The S3 Lifecycle policy transitions the data to Glacier Deep Archive after 1 year for long-term archival. S3 Object Lock in compliance mode prevents any user from deleting or overwriting objects for the specified retention period. Glacier Deep Archive provides very high durability and the lowest storage cost for long-term archival. Compliance mode ensures no one can override or change the retention settings even if policies change. This meets all the requirements - immediate access for 1 year, archived for 9 years, unable to delete for 10 years, maximum resiliency

upvoted 10 times

satyaamm Most Recent 5 months, 1 week ago

Selected Answer: C

S3 lifecycle policies are required to transition from immediate availability to archived form. Also the S3 object lock helps prevent deletion of the bucket.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: C

Ans C - S3 Glacier after year 1 in compliance mode with object lock (=immutable lock)

upvoted 2 times

Ruffyt 1 year, 7 months ago

No one at the company, including administrative users and root users, can be able to delete the records during the entire 10-year period = Compliance Mode

upvoted 5 times

axelrodb 1 year, 9 months ago

Selected Answer: C

To meet the requirements of immediately accessible records for 1 year and then archived for an additional 9 years with maximum resiliency, we can use S3 Lifecycle policy to transition records from S3 Standard to S3 Glacier Deep Archive after 1 year. And to ensure that the records cannot be deleted by anyone, including administrative and root users, we can use S3 Object Lock in compliance mode for a period of 10 years. Therefore, the correct answer is option C.

Reference: <https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.htmls>

upvoted 5 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

No one at the company, including administrative users and root users, can be able to delete the records during the entire 10-year period = Compliance Mode

upvoted 2 times

miki111 1 year, 10 months ago

Option C is the correct answer

upvoted 2 times

MutiverseAgent 1 year, 11 months ago

Why not A? Move all files to S3 Glacier instant retrieval (Cheaper than S3) and then move files older than a year to S3 Deep archive.

upvoted 1 times

dhax12 1 year, 8 months ago

Put entire 10 years to Glacier means it's not accessible for the 1 year window. Hence wrong answer.

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: C

To prevent deletion of records during the entire 10-year period, you can utilize S3 Object Lock feature. By enabling it in compliance mode, you can set a retention period on the objects, preventing any user, including administrative and root users, from deleting records.

A: S3 Glacier is suitable for long-term archival, it may not provide immediate accessibility for the first year as required.

B: Intelligent-Tiering may not offer the most cost-effective archival storage option for extended 9-year period. Changing the IAM policy after 10 years to allow deletion also introduces manual steps and potential human error.

D: While S3 One Zone-IA can provide cost savings, it doesn't offer the same level of resiliency as S3 Glacier Deep Archive for long-term archival.

upvoted 4 times

11pantheman11 2 years, 1 month ago

Selected Answer: C

In compliance mode, a protected object version can't be overwritten or deleted by any user, including the root user in your AWS account.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock-overview.html>

upvoted 4 times

athiha 2 years, 3 months ago

Selected Answer: C

Retention Period: A period is specified by Days & Years.

With Retention Compliance Mode, you can't change/adjust (even by the account root user) the retention mode during the retention period while all objects within the bucket are Locked.

With Retention Governance mode, a less restrictive mode, you can grant special permission to a group of users to adjust the Lock settings by using S3:BypassGovernanceRetention.

Legal Hold: It's On/Off setting on an object version. There is no retention period. If you enable Legal Hold on specific object version, you will not be able to delete or override that specific object version. It needs S:PutObjectLegalHold as a permission.

upvoted 5 times

WheracanIstart 2 years, 3 months ago

Selected Answer: C

S3 Glacier Deep Archive all day....

upvoted 1 times

SilentMilli 2 years, 5 months ago

Selected Answer: C

Use an S3 Lifecycle policy to transition the records from S3 Standard to S3 Glacier Deep Archive after 1 year. Use S3 Object Lock in compliance mode for a period of 10 years.

upvoted 2 times

k1kavi1 2 years, 5 months ago

Selected Answer: C

Use S3 Object Lock in compliance mode

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock-overview.html>

upvoted 4 times

pazabal 2 years, 5 months ago

Selected Answer: C

C, A lifecycle set to transition from standard to Glacier deep archive and use lock for the delete requirement

A, B and D don't meet the requirements

upvoted 2 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: C

C. Use an S3 Lifecycle policy to transition the records from S3 Standard to S3 Glacier Deep Archive after 1 year. Use S3 Object Lock in compliance mode for a period of 10 years.

To meet the requirements, the company could use an S3 Lifecycle policy to transition the records from S3 Standard to S3 Glacier Deep Archive after 1 year. S3 Glacier Deep Archive is Amazon's lowest-cost storage class, specifically designed for long-term retention of data that is accessed rarely. This would allow the company to store the records with maximum resiliency and at the lowest possible cost.

upvoted 4 times

Buruguduystunstugudunstuy 2 years, 5 months ago

To ensure that the records are not deleted during the entire 10-year period, the company could use S3 Object Lock in compliance mode. S3 Object Lock allows the company to apply a retention period to objects in S3, preventing the objects from being deleted until the retention period expires. By using S3 Object Lock in compliance mode, the company can ensure that the records are not deleted by anyone, including administrative users and root users, during the entire 10-year period.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #54


Topic 1

A company runs multiple Windows workloads on AWS. The company's employees use Windows file shares that are hosted on two Amazon EC2 instances. The file shares synchronize data between themselves and maintain duplicate copies. The company wants a highly available and durable storage solution that preserves how users currently access the files. What should a solutions architect do to meet these requirements?

- A. Migrate all the data to Amazon S3. Set up IAM authentication for users to access files.
- B. Set up an Amazon S3 File Gateway. Mount the S3 File Gateway on the existing EC2 instances.
- C. Extend the file share environment to Amazon FSx for Windows File Server with a Multi-AZ configuration. Migrate all the data to FSx for Windows File Server. **Most Voted**
- D. Extend the file share environment to Amazon Elastic File System (Amazon EFS) with a Multi-AZ configuration. Migrate all the data to Amazon EFS.

Correct Answer: C

Community vote distribution



Comments

k1kavi1 **Highly Voted** 2 years, 5 months ago

Selected Answer: C

EFS is not supported on Windows instances

<https://docs.aws.amazon.com/AWSEC2/latest/WindowsGuide/AmazonEFS.html>

Amazon FSx for Windows File Server provides fully managed Microsoft Windows file servers, backed by a fully native Windows file system.

<https://docs.aws.amazon.com/fsx/latest/WindowsGuide/what-is.html>

upvoted 21 times

Buruguduystunstugudunstuy **Highly Voted** 2 years, 6 months ago

Selected Answer: C

Windows file shares = Amazon FSx for Windows File Server

Hence, the correct answer is C

upvoted 8 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Taking back this answer. As explained in the latest update.

CORRECT

D: Extend the file share environment to Amazon Elastic File System (Amazon EFS) with a Multi-AZ configuration. Migrate all the data to Amazon EFS.

upvoted 1 times

pentium75 1 year, 5 months ago

No, users should continue using SMB, which EFS doesn't support

upvoted 7 times

satyaamm Most Recent 5 months, 1 week ago

Selected Answer: C

FSx is used for windows file systems whereas EFS is used for Linux file systems.

upvoted 2 times

PaulGa 9 months ago

Selected Answer: C

Ans C - use FSx for Windows File Server - why create extra work by migrating to EFS? Add Multi-AZ configuration for availability

upvoted 2 times

MatAlves 10 months ago

Amazon EFS (Linux) and Amazon FSx (Windows) provide highly durable and available file systems that can span multiple Availability Zones. Both solutions are designed to deliver high performance, however, when choosing to use network file systems consider the access patterns.

<https://docs.aws.amazon.com/wellarchitected/latest/sap-lens/best-practice-14-3.html>

upvoted 2 times

SMALLE 1 year, 4 months ago

Selected Answer: C

FSx=Windows

<https://docs.aws.amazon.com/fsx/latest/WindowsGuide/what-is.html>

upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

Windows workload rules out S3 and EFS as they cannot be mounted directly on Windows. S3 File Gateway is mainly for on-prem to AWS which is not a requirement here as company is already in AWS.

C meets all the requirements.

upvoted 4 times

jithlathi 1 year, 5 months ago

Selected Answer: C

With Amazon FSx for Windows File Server, you can enjoy a native Windows file server experience with a fully managed, scalable, and highly dependable file storage solution

upvoted 3 times

Ruffyt 1 year, 7 months ago

<https://aws.amazon.com/fsx/windows/faqs/>

Thousands of compute instances and devices can access a file system concurrently.

upvoted 2 times

AWSSStudyBuddy 1 year, 7 months ago

Selected Answer: C

With Amazon FSx for Windows File Server, you can enjoy a native Windows file server experience with a fully managed, scalable, and highly dependable file storage solution. Rich administrative features including end-user file recovery, user quotas,

scalable, and highly dependable file storage solution. Rich administrative features including end-user file recovery, user quotas, and Microsoft Active Directory integration are all provided by this Windows Server-based system.

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

The key reasons are:

FSx for Windows provides fully managed Windows-native SMB file shares that are accessible from Windows clients.

It allows seamlessly migrating the existing Windows file shares to FSx shares without disrupting users.

The Multi-AZ configuration provides high availability and durability for file storage.

Users can continue to access files the same way over SMB without any changes.

It is optimized for Windows workloads and provides features like user quotas, ACLs, AD integration.

Data is stored on SSDs with automatic backups for resilience.

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

The company wants a highly available and durable storage solution that preserves how users currently access the files = Amazon FSx for Windows File Server

upvoted 2 times

miki111 1 year, 10 months ago

Option C is the correct answer

upvoted 1 times

cookieMr 1 year, 11 months ago

Selected Answer: C

Migrating all the data to FSx for Windows File Server allows you to preserve existing user access method and maintain compatibility with Windows file shares. Users can continue accessing files using the same method as before, without any disruptions.

A: S3 is a highly durable object storage service, it is not designed to directly host Windows file shares. Implementing IAM authentication for file access would require significant changes to existing user access method.

B: S3 File Gateway can provide access to Amazon S3 objects through standard file protocols, it may not be ideal solution for preserving existing user access method and maintaining Windows file shares.

D: Although Amazon EFS provides highly available and durable file storage, it may not directly support the existing Windows file shares and their access method.

upvoted 5 times

11pantheman11 2 years, 1 month ago

Selected Answer: C

<https://aws.amazon.com/fsx/windows/faqs/>

Thousands of compute instances and devices can access a file system concurrently.

EFS does not support Windows

upvoted 3 times

cheese929 2 years, 1 month ago

Selected Answer: C

C is correct. Amazon FSx for Windows File Server.

upvoted 3 times

satosis 2 years, 1 month ago

Selected Answer: C

EFS is not supported on Windows instances

<https://docs.aws.amazon.com/AWSEC2/latest/WindowsGuide/AmazonEFS.html>

Amazon FSx for Windows File Server provides fully managed Microsoft Windows file servers, backed by a fully native Windows file system.

<https://docs.aws.amazon.com/fsx/latest/WindowsGuide/what-is.html>

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #55

Topic 1

A solutions architect is developing a VPC architecture that includes multiple subnets. The architecture will host applications that use Amazon EC2 instances and Amazon RDS DB instances. The architecture consists of six subnets in two Availability Zones. Each Availability Zone includes a public subnet, a private subnet, and a dedicated subnet for databases. Only EC2 instances that run in the private subnets can have access to the RDS databases.

Which solution will meet these requirements?

- A. Create a new route table that excludes the route to the public subnets' CIDR blocks. Associate the route table with the database subnets.
- B. Create a security group that denies inbound traffic from the security group that is assigned to instances in the public subnets. Attach the security group to the DB instances.
- C. Create a security group that allows inbound traffic from the security group that is assigned to instances in the private subnets. Attach the security group to the DB instances. **Most Voted**
- D. Create a new peering connection between the public subnets and the private subnets. Create a different peering connection between the private subnets and the database subnets.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Sinaneos **Highly Voted** 2 years, 8 months ago

Selected Answer: C

A: doesn't fully configure the traffic flow
 B: security groups don't have deny rules
 D: peering is mostly between VPCs, doesn't really help here

answer is C, most mainstream way
 upvoted 55 times

Gary_Phillips_2007 **Highly Voted** 2 years, 3 months ago

Just took the exam today and EVERY ONE of the questions came from this dump. Memorize it all. Good luck.
 upvoted 36 times

orhan64 1 year, 10 months ago

Hey bro, did you buy premium access?

upvoted 5 times

PaulGa Most Recent 9 months ago

Selected Answer: C

Ans C - only one that makes sense: "...a security group that allows inbound traffic from the security group that is assigned to instances in the private subnets" - operative word: "allows"

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

A: route table that connect... no idea what this option is trying to do but won't work for RDS

B: SG are deny by default

D: Peering connection between subnets? No idea what this is but happy to learn if such a thing exists.

C: SG to allow input to private subnet means everything else will be blocked. Attaching this SG to DB instance means it will block everything except the private subnet instances which is where the required EC2 instances are.

upvoted 5 times

AWSStudyBuddy 1 year, 7 months ago

Selected Answer: C

RDS databases can only be accessed by EC2 instances located in private subnets: From the security group given to instances in the private subnets, the DB instances' security group will permit incoming traffic. Because of this, the RDS databases will only be accessible by EC2 instances located on the private subnets.

Because of its safe architecture, Every other source of incoming traffic will be blocked by the security group that is linked to the database instances. The RDS databases will be better shielded from unwanted access thanks to this.

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

The key reasons are:

Using security groups to control access between resources is a standard practice in VPCs.

The security group attached to the RDS DB instances can allow inbound traffic from the security group for the EC2 instances in the private subnets.

This allows only those EC2 instances in the private subnets to connect to the databases, meeting the requirements.

Route tables, peering connections, and denying public subnet access would not achieve the needed selectivity of allowing only the private subnet EC2 instances.

Security groups provide stateful filtering at the instance level for precise access control.

upvoted 3 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

Security groups only have allow rules

upvoted 3 times

praveenvky83 1 year, 10 months ago

Selected Answer: C

optoin C

upvoted 1 times

miki111 1 year, 10 months ago

Option C is the correct answer

upvoted 1 times

cookieMr 1 year, 11 months ago

Selected Answer: C

Creating security group that allows inbound traffic from security group assigned to instances in private subnets ensures that

creating security group that allows inbound traffic from security group assigned to instances in private subnets ensures that only EC2 running in private subnets can access the RDS databases. By associating security group with DB, you restrict access to only instances that belong to designated security group.

A: This approach may help control routing within VPC, it does not address the specific access requirement between EC2 instances and RDS databases.

B: Using a deny rule in a security group can lead to complexities and potential misconfigurations. It is generally recommended to use allow rules to explicitly define access permissions.

D: Peering connections enable communication between different VPCs or VPCs in different regions, and they are not necessary for restricting access between subnets within the same VPC.

upvoted 5 times

Bmarodi 2 years ago

Selected Answer: C

Option C meets the requirements.

upvoted 1 times

Abrar2022 2 years ago

By default, a security group is set up with rules that deny all inbound traffic and permit all outbound traffic.

upvoted 2 times

water314 2 years, 1 month ago

Selected Answer: C

CCCCCCCCCCC

upvoted 1 times

SilentMilli 2 years, 5 months ago

Selected Answer: C

Create a security group that allows inbound traffic from the security group that is assigned to instances in the private subnets. Attach the security group to the DB instances. This will allow the EC2 instances in the private subnets to have access to the RDS databases while denying access to the EC2 instances in the public subnets.

upvoted 2 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: C

The solution that meets the requirements described in the question is option C: Create a security group that allows inbound traffic from the security group that is assigned to instances in the private subnets. Attach the security group to the DB instances.

In this solution, the security group applied to the DB instances allows inbound traffic from the security group assigned to instances in the private subnets. This ensures that only EC2 instances running in the private subnets can have access to the RDS databases.

upvoted 3 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option A, creating a new route table that excludes the route to the public subnets' CIDR blocks and associating it with the database subnets, would not meet the requirements because it would block all traffic to the database subnets, not just traffic from the public subnets.

Option B, creating a security group that denies inbound traffic from the security group assigned to instances in the public subnets and attaching it to the DB instances, would not meet the requirements because it would allow all traffic from the private subnets to reach the DB instances, not just traffic from the security group assigned to instances in the private subnets.

Option D, creating a new peering connection between the public subnets and the private subnets and a different peering connection between the private subnets and the database subnets, would not meet the requirements because it would allow all traffic from the private subnets to reach the DB instances, not just traffic from the security group assigned to instances in the private subnets.

upvoted 1 times

Nandan747 2 years, 5 months ago

Selected Answer: C

The real trick is between B and C. A and D are ruled out for obvious reasons.

B is wrong as you cannot have deny type rules in Security groups.
So- C is the right answer.

upvoted 4 times

ashish_t 2 years, 6 months ago

Selected Answer: C

The key is "Only EC2 instances that run in the private subnets can have access to the RDS databases"
The answer is C.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #56

Topic 1


A company has registered its domain name with Amazon Route 53. The company uses Amazon API Gateway in the ca-central-1 Region as a public interface for its backend microservice APIs. Third-party services consume the APIs securely. The company wants to design its API Gateway URL with the company's domain name and corresponding certificate so that the third-party services can use HTTPS.

Which solution will meet these requirements?

- A. Create stage variables in API Gateway with Name="Endpoint-URL" and Value="Company Domain Name" to overwrite the default URL. Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM).
- B. Create Route 53 DNS records with the company's domain name. Point the alias record to the Regional API Gateway stage endpoint. Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the us-east-1 Region.
- C. Create a Regional API Gateway endpoint. Associate the API Gateway endpoint with the company's domain name. Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the same Region. Attach the certificate to the API Gateway endpoint. Configure Route 53 to route traffic to the API Gateway endpoint. **Most Voted**
- D. Create a Regional API Gateway endpoint. Associate the API Gateway endpoint with the company's domain name. Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the us-east-1 Region. Attach the certificate to the API Gateway APIs. Create Route 53 DNS records with the company's domain name. Point an A record to the company's domain name.

Correct Answer: C

Community vote distribution



Comments

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: C

The correct solution to meet these requirements is option C.

To design the API Gateway URL with the company's domain name and corresponding certificate, the company needs to do the

following:

1. Create a Regional API Gateway endpoint: This will allow the company to create an endpoint that is specific to a region.
2. Associate the API Gateway endpoint with the company's domain name: This will allow the company to use its own domain name for the API Gateway URL.
3. Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the same Region: This will allow the company to use HTTPS for secure communication with its APIs.
4. Attach the certificate to the API Gateway endpoint: This will allow the company to use the certificate for securing the API Gateway URL.
5. Configure Route 53 to route traffic to the API Gateway endpoint: This will allow the company to use Route 53 to route traffic to the API Gateway URL using the company's domain name.

upvoted 66 times

t0nx 1 year, 6 months ago

Why the "reveal solution" most of the time gives the wrong answer ?

upvoted 13 times

yaralexie 4 months, 2 weeks ago

☐ happens to me too

upvoted 1 times

wharftargo 1 year, 4 months ago

i read this before that they can't give 100% of the right answers legally or something

upvoted 9 times

aadityaravi8 1 year, 11 months ago

google bard reply..

upvoted 4 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option C includes all the necessary steps to meet the requirements, hence it is the correct solution.

Options A and D do not include the necessary steps to associate the API Gateway endpoint with the company's domain name and attach the certificate to the endpoint.

Option B includes the necessary steps to associate the API Gateway endpoint with the company's domain name and attach the certificate, but it imports the certificate into the us-east-1 Region instead of the ca-central-1 Region where the API Gateway is located.

upvoted 12 times

masetromain Highly Voted 2 years, 8 months ago

Selected Answer: C

I think the answer is C. we don't need to attach a certificate in us-east-1, if is not for cloudfront. In our case the target is ca-central-1.

upvoted 33 times

Valero_ 2 years, 8 months ago

I think that is C too, the target would be the same Region.

<https://docs.aws.amazon.com/apigateway/latest/developerguide/apigateway-regional-api-custom-domain-create.html>

upvoted 8 times

MutiverseAgent 1 year, 11 months ago

Agree, C is correct by using the API Gateway option "Custom domain names"

<https://docs.aws.amazon.com/apigateway/latest/developerguide/how-to-custom-domains.html>

upvoted 2 times

FlyingHawk Most Recent 4 months, 3 weeks ago

Selected Answer: C

I selected B initially as I mixed API Gateway with CloudFront, for CloudFront, it is a global service, the cert needs to be in us-east-1 region as CF's control plane is in us-east-1 region, but for regional API Gateway, ELB, they should be in the same region as the service.

<https://docs.aws.amazon.com/acm/latest/userguide/import-certificate.html>

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/cnames-and-https-requirements.html>

upvoted 1 times

PaulGa 9 months ago

Selected Answer: C

Ans C - as per Buruguduystunstugudunstuy (1 year, 8 months ago) for the reasons therein...

Not sure what Ans D is addressing...

upvoted 2 times

jatric 11 months, 2 weeks ago

Selected Answer: C

Option C has all the steps to meet the requirement and attach certificate in the same region

upvoted 2 times

ChymKuBoy 12 months ago

Selected Answer: C

C for sure

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

BD are wrong because they are in wrong regions.

A. Does not help with R53 routing to API Gateway and not sure what it's trying to do here

C is correct

upvoted 2 times

bujuman 1 year, 5 months ago

Selected Answer: C

Important

For an API Gateway Regional custom domain name, you must request or import the certificate in the same Region as your API.

upvoted 4 times

debasishdtta 1 year, 5 months ago

Selected Answer: D

All certificates in ACM are regional resources, including the certificates that you import. To use the same certificate with Elastic Load Balancing load balancers in different AWS Regions, you must import the certificate into each Region where you want to use it. To use a certificate with Amazon CloudFront, you must import it into the US East (N. Virginia) Region.

<https://docs.aws.amazon.com/acm/latest/userguide/import-certificate.html>

upvoted 3 times

EtherealBagel 1 year, 6 months ago

Only if the API Gateway is global then the corresponding AWS ACM Certificate must be placed in us-east-1

upvoted 1 times

luongtrann 1 year, 7 months ago

Selected Answer: C

Correct answer

upvoted 1 times

Abitek007 1 year, 8 months ago

Selected Answer: D

A records support Elasticity and load balancing and by default resilience is Key in any configuration in AWS

upvoted 2 times

Abitek007 1 year, 8 months ago

now I am confused, I would have chosen C, but with a Closer look D might be right, because of the A records and again the region used and not stated can be for resilience. I think? can someone clarify

upvoted 2 times

OctavioBatera 1 year, 2 months ago

I think C is the correct answer, because the DNS record in this case must be an alias (cname). DNS A record is for IP address. Here some documentation that can be useful:

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-to-api-gateway.html>

upvoted 3 times

paniya93 1 year, 8 months ago

Selected Answer: C

Explain why this saying a different region which not mentioned in the Q.

upvoted 1 times

Hassao0 1 year, 9 months ago

c is right

The other options have various issues:

Option A: Using stage variables and importing certificates into ACM is not sufficient for achieving the requirement of associating a custom domain and certificate with the API Gateway endpoint.

Option B: While it mentions importing the certificate into ACM, it doesn't address the need for a Regional API Gateway or the appropriate region for the certificate.

Option D: Using certificates from the us-east-1 region for a Regional API Gateway might cause issues. Additionally, it doesn't provide clear details on how to associate the domain name and certificate with the API Gateway endpoint.

upvoted 3 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

C is the correct solution.

To use a custom domain name with HTTPS for API Gateway:

The API Gateway endpoint needs to be Regional, not private or edge-optimized.

The ACM certificate must be requested in the same region as the API Gateway endpoint.

The custom domain name is then mapped to the Regional API endpoint under API Gateway domain names.

Route 53 is configured to route traffic to the API Gateway regional domain.

The ACM certificate is attached to the API Gateway domain name to enable HTTP

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the same Region.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #57

Topic 1

A company is running a popular social media website. The website gives users the ability to upload images to share with other users. The company wants to make sure that the images do not contain inappropriate content. The company needs a solution that minimizes development effort.

What should a solutions architect do to meet these requirements?

- A. Use Amazon Comprehend to detect inappropriate content. Use human review for low-confidence predictions.
- B. Use Amazon Rekognition to detect inappropriate content. Use human review for low-confidence predictions. **Most Voted**
- C. Use Amazon SageMaker to detect inappropriate content. Use ground truth to label low-confidence predictions.
- D. Use AWS Fargate to deploy a custom machine learning model to detect inappropriate content. Use ground truth to label low-confidence predictions.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: B

The best solution to meet these requirements would be option B: Use Amazon Rekognition to detect inappropriate content, and use human review for low-confidence predictions.

Amazon Rekognition is a cloud-based image and video analysis service that can detect inappropriate content in images using its pre-trained label detection model. It can identify a wide range of inappropriate content, including explicit or suggestive adult content, violent content, and offensive language. The service provides high accuracy and low latency, making it a good choice for this use case.

upvoted 25 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option A, using Amazon Comprehend, is not a good fit for this use case because Amazon Comprehend is a natural language processing service that is designed to analyze text, not images.

Option C, using Amazon SageMaker to detect inappropriate content, would require significant development effort to build and train a custom machine learning model. It would also require a large dataset of labeled images to train the model, which may be time-consuming and expensive to obtain.

Option D, using AWS Fargate to deploy a custom machine learning model, would also require significant development effort and a large dataset of labeled images. It may not be the most efficient or cost-effective solution for this use case.

In summary, the best solution is to use Amazon Rekognition to detect inappropriate content in images, and use human review for low-confidence predictions to ensure that all inappropriate content is detected.

upvoted 16 times

masetromain Highly Voted 2 years, 8 months ago

Selected Answer: B

Good Answer is B :

<https://docs.aws.amazon.com/rekognition/latest/dg/moderation.html?pg=ln&sec=ft>

upvoted 14 times

satyaamm Most Recent 5 months, 1 week ago

Selected Answer: B

Amazon Rekognition is a ML service by AWS for image and video analysis so it is the most suitable option here.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: B

Ans B - I'll go with ans B reluctantly because Rekognition seems to be primarily aimed at video/image content as opposed to language/text... and only because you can add the latter "You can add features [to Rekognition] that detect objects, text, unsafe content, analyze images/videos, and compare faces to your application using Rekognition's APIs."

(<https://docs.aws.amazon.com/rekognition/latest/dg/what-is.html>)

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

<https://aws.amazon.com/rekognition/>

Automate and lower the cost of your image recognition and video analysis with machine learning

upvoted 2 times

awsgeek75 1 year, 5 months ago

<https://aws.amazon.com/rekognition/content-moderation/>

Amazon Rekognition Content Moderation automates and streamlines your image and video moderation workflows using machine learning (ML), without requiring ML experience.

upvoted 2 times

slimen 1 year, 7 months ago

Selected Answer: B

comprehend is for NLP

sagemaker is for training and deploying ML and AI models

deploying custom models using fargate requires time and development effort which is not recommended by the question

upvoted 4 times

Ruffyit 1 year, 7 months ago

<https://docs.aws.amazon.com/rekognition/latest/dg/moderation.html?pg=ln&sec=ft>

upvoted 2 times

AWSStudyBuddy 1 year, 7 months ago

Selected Answer: B

You can easily incorporate image and video analysis to your applications with the help of Amazon Rekognition. Numerous functions are available to it, including as facial analysis, image classification, and object and scene identification.

DetectModerationLabels is an operation that may be used with Amazon Rekognition to identify incorrect content in photos. By using this procedure, photos with violent, drug-related, tobacco-related, alcohol-related, hate-filled, or provocative material can be identified.

upvoted 3 times

Svruis 1 year, 10 months ago

Selected Answer: B

B is the best solution as far

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

Amazon Rekognition is a fully managed service that provides image and video analysis capabilities. It can be used to detect inappropriate content in images, such as nudity, violence, and hate speech.

Amazon Rekognition is a good choice for this solution because it is a managed service, which means that the company does not have to worry about managing the infrastructure or the machine learning model. Rekognition is also highly accurate, and it can be used to detect a wide range of inappropriate content

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: B

Amazon Rekognition to the rescue...whoosh!

upvoted 1 times

cookieMr 1 year, 11 months ago

Using Amazon Rekognition for content moderation is a cost-effective and efficient solution that reduces the need for developing and training custom machine learning models, making it the best option in terms of minimizing development effort.

A. Amazon Comprehend is a natural language processing service provided by AWS, primarily focused on text analysis rather than image analysis.

C. Amazon SageMaker is a comprehensive machine learning service that allows you to build, train, and deploy custom machine learning models. It requires significant development effort to build and train a custom model. In addition, utilizing ground truth to label low-confidence predictions would further add to the development complexity and maintenance overhead.

D. Similar to C, using AWS Fargate to deploy a custom machine learning model requires significant development effort.

upvoted 3 times

krajar 2 years, 2 months ago

Selected Answer: B

Amazon Rekognition is a cloud-based image and video analysis service that can detect inappropriate content in images using its pre-trained label detection model. It can identify a wide range of inappropriate content, including explicit or suggestive adult content, violent content, and offensive language.

upvoted 2 times

career360guru 2 years, 5 months ago

Selected Answer: B

Option B

upvoted 1 times

Shasha1 2 years, 6 months ago

B
AWS Rekognition to detect inappropriate content and use human review for low-confidence predictions. This option minimizes development effort because Amazon Rekognition is a pre-built machine learning service that can detect inappropriate content. Using human review for low-confidence predictions allows for more accurate detection of inappropriate content.

upvoted 2 times

Wpcorgan 2 years, 6 months ago

B is correct

upvoted 1 times

ArielSchivo 2 years, 7 months ago

Selected Answer: B

Option B.

<https://docs.aws.amazon.com/rekognition/latest/dg/a2i-rekognition.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #58

Topic 1

A company wants to run its critical applications in containers to meet requirements for scalability and availability. The company prefers to focus on maintenance of the critical applications. The company does not want to be responsible for provisioning and managing the underlying infrastructure that runs the containerized workload.

What should a solutions architect do to meet these requirements?

- A. Use Amazon EC2 instances, and install Docker on the instances.
- B. Use Amazon Elastic Container Service (Amazon ECS) on Amazon EC2 worker nodes.
- C. Use Amazon Elastic Container Service (Amazon ECS) on AWS Fargate. **Most Voted**
- D. Use Amazon EC2 instances from an Amazon Elastic Container Service (Amazon ECS)-optimized Amazon Machine Image (AMI).

Correct Answer: C*Community vote distribution*

C (100%)

Comments**masetromain** **Highly Voted** 2 years, 8 months ago**Selected Answer: C**

Good answer is C:

AWS Fargate is a serverless, pay-as-you-go compute engine that lets you focus on building applications without having to manage servers. AWS Fargate is compatible with Amazon Elastic Container Service (ECS) and Amazon Elastic Kubernetes Service (EKS).

<https://aws.amazon.com/fr/fargate/>

upvoted 28 times

cookieMr **Highly Voted** 1 year, 11 months ago**Selected Answer: C**

Using ECS on Fargate allows you to run containers without the need to manage the underlying infrastructure. Fargate abstracts away the underlying EC2 and provides serverless compute for containers.

A. This option would require manual provisioning and management of EC2, as well as installing and configuring Docker on those instances. It would introduce additional overhead and responsibilities for maintaining the underlying infrastructure.

B. While this option leverages ECS to manage containers, it still requires provisioning and managing EC2 to serve as worker nodes. It adds complexity and maintenance overhead compared to the serverless nature of Fargate.

D. This option still involves managing and provisioning EC2, even though an ECS-optimized AMI simplifies the process of setting up EC2 for running ECS. It does not provide the level of serverless abstraction and ease of management offered by Fargate.

upvoted 8 times

PaulGa Most Recent 9 months ago

Selected Answer: C

Ans C - use Fargate to do all the management / deployment (which the company doesn't want to do)

upvoted 2 times

Hkayne 1 year, 1 month ago

Selected Answer: C

ECS FARGATE

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

Managed containers = Fargate

upvoted 4 times

Ruffyt 1 year, 7 months ago

AWS Fargate is a serverless, pay-as-you-go compute engine that lets you focus on building applications without having to manage servers. AWS Fargate is compatible with Amazon Elastic Container Service (ECS) and Amazon Elastic Kubernetes Service (EKS).

upvoted 2 times

AWSStudyBuddy 1 year, 7 months ago

Selected Answer: C

In order to execute containerized apps without having to manage servers, AWS Fargate is a serverless compute engine for Amazon ECS. Amazon Elastic Compute Cloud (Amazon EC2) instance clusters no longer require provisioning, configuring, or scaling thanks to AWS Fargate. So that you can concentrate on developing and maintaining your applications, AWS Fargate handles the monotonous, repetitive labor of managing servers.

upvoted 2 times

Teruteru 1 year, 9 months ago

Option C is the correct answer.

upvoted 1 times

Syruis 1 year, 10 months ago

Selected Answer: C

C for Fargate

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

The company does not want to be responsible for provisioning and managing the underlying infrastructure that runs the containerized workload = Serverless compute for containers = AWS Fargate

upvoted 2 times

miki111 1 year, 10 months ago

Option C is the correct answer

upvoted 1 times

cheese929 2 years, 1 month ago

Selected Answer: C

AWS Fargate is a technology that you can use with Amazon ECS to run containers without having to manage servers or clusters.

AWS Fargate is a technology that you can use with Amazon ECS to run containers without having to manage servers or clusters of Amazon EC2 instances.

<https://docs.aws.amazon.com/AmazonECS/latest/userguide/what-is-fargate.html>

upvoted 2 times

SilentMilli 2 years, 5 months ago

Selected Answer: C

ECS + Fargate

upvoted 3 times

gustavtd 2 years, 5 months ago

Selected Answer: C

AWS Fargate will hide all the complexity for you

upvoted 1 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: C

C. Use Amazon Elastic Container Service (Amazon ECS) on AWS Fargate.

AWS Fargate is a fully managed container execution environment that runs containers without the need to provision and manage underlying infrastructure. This makes it a good choice for companies that want to focus on maintaining their critical applications and do not want to be responsible for provisioning and managing the underlying infrastructure.

Option A involves installing Docker on Amazon EC2 instances, which would still require the company to manage the underlying infrastructure. Option B involves using Amazon ECS on Amazon EC2 worker nodes, which would also require the company to manage the underlying infrastructure. Option D involves using Amazon EC2 instances from an Amazon ECS-optimized Amazon Machine Image (AMI), which would also require the company to manage the underlying infrastructure.

upvoted 3 times

career360guru 2 years, 5 months ago

Selected Answer: C

Option C

upvoted 1 times

benaws 2 years, 6 months ago

Selected Answer: C

Obviously anything with EC2 in the answer is wrong...

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #59

Topic 1


A company hosts more than 300 global websites and applications. The company requires a platform to analyze more than 30 TB of clickstream data each day.

What should a solutions architect do to transmit and process the clickstream data?

- A. Design an AWS Data Pipeline to archive the data to an Amazon S3 bucket and run an Amazon EMR cluster with the data to generate analytics.
- B. Create an Auto Scaling group of Amazon EC2 instances to process the data and send it to an Amazon S3 data lake for Amazon Redshift to use for analysis.
- C. Cache the data to Amazon CloudFront. Store the data in an Amazon S3 bucket. When an object is added to the S3 bucket, run an AWS Lambda function to process the data for analysis.
- D. Collect the data from Amazon Kinesis Data Streams. Use Amazon Kinesis Data Firehose to transmit the data to an Amazon S3 data lake. Load the data in Amazon Redshift for analysis. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: D

Option D is the most appropriate solution for transmitting and processing the clickstream data in this scenario.

Amazon Kinesis Data Streams is a highly scalable and durable service that enables real-time processing of streaming data at a high volume and high rate. You can use Kinesis Data Streams to collect and process the clickstream data in real-time.

Amazon Kinesis Data Firehose is a fully managed service that loads streaming data into data stores and analytics tools. You can use Kinesis Data Firehose to transmit the data from Kinesis Data Streams to an Amazon S3 data lake.

Once the data is in the data lake, you can use Amazon Redshift to load the data and perform analysis on it. Amazon Redshift is a fully managed, petabyte-scale data warehouse service that allows you to quickly and efficiently analyze data using SQL and your existing business intelligence tools.

upvoted 40 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option A, which involves using AWS Data Pipeline to archive the data to an Amazon S3 bucket and running an Amazon EMR cluster with the data to generate analytics, is not the most appropriate solution because it does not involve real-time processing of the data.

Option B, which involves creating an Auto Scaling group of Amazon EC2 instances to process the data and sending it to an Amazon S3 data lake for Amazon Redshift to use for analysis, is not the most appropriate solution because it does not involve a fully managed service for transmitting the data from the processing layer to the data lake.

Option C, which involves caching the data to Amazon CloudFront, storing the data in an Amazon S3 bucket, and running an AWS Lambda function to process the data for analysis when an object is added to the S3 bucket, is not the most appropriate solution because it does not involve a scalable and durable service for collecting and processing the data in real-time.

upvoted 10 times

MutiverseAgent 1 year, 11 months ago

The question does not say that real-time is needed here

upvoted 3 times

pentium75 1 year, 5 months ago

Question asks how to "transmit and process the clickstream data", NOT how to analyze it. Thus D.

upvoted 2 times

ArielSchivo Highly Voted 2 years, 7 months ago

Selected Answer: D

Option D.

<https://aws.amazon.com/es/blogs/big-data/real-time-analytics-with-amazon-redshift-streaming-ingestion/>

upvoted 17 times

RBSK 2 years, 6 months ago

Unsure if this is right URL for this scenario. Option D is referring to S3 and then Redshift. Whereas URL discuss about eliminating S3 :- We're excited to launch Amazon Redshift streaming ingestion for Amazon Kinesis Data Streams, which enables you to ingest data directly from the Kinesis data stream without having to stage the data in Amazon Simple Storage Service (Amazon S3). Streaming ingestion allows you to achieve low latency in the order of seconds while ingesting hundreds of megabytes of data into your Amazon Redshift cluster.

upvoted 5 times

satyaamm Most Recent 5 months, 1 week ago

Selected Answer: D

Kinesis data firehouse is the most suitable for streaming data and Redshift is the most suitable for large data sets.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: D

Ans D - using Kinesis Streams / Firehouse (data in/out) is fast and reliable. Using Redshift allows all sorts of permutations of data analyses and interfacing to user apps

upvoted 2 times

effiecanode 11 months, 1 week ago

D is the best option

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

A: Not sure how recent this question is but Data Pipeline is not really a product AWS is recommending anymore
<https://docs.aws.amazon.com/datapipeline/latest/DeveloperGuide/what-is-datapipeline.html>

B: 30TB of clickstream data could be done with EC2 but it would be challenging

C: CloudFront is for CDN and caching and mostly outgoing data, not incoming.

D: Kinesis, S3 data lake and Redshift will work perfectly for this case

upvoted 4 times

upvoted 4 times

clumsyninja4life 1 year, 5 months ago

Selected Answer: A

The answer should be A. Clickstream does not mean real time, it just means they capture user interactions on the web page. Kinesis data streaming is not required. Furthermore, redshift is a data warehousing solution, it cant run complex analysis as well as EMR. My vote goes for A

upvoted 1 times

pentium75 1 year, 5 months ago

Question asks how to "transmit and process the clickstream data", NOT how to analyze it. Also question does NOT ask how to archive the data (as is mentioned in A). Thus D.

upvoted 1 times

Reckless_Jas 1 year, 10 months ago

when you see clickstream data, think about Kinesis Data Stream

upvoted 6 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: D

The key reasons are:

Kinesis Data Streams can continuously capture and ingest high volumes of clickstream data in real-time. This handles the large 30TB daily data intake.

Kinesis Firehose can automatically load the streaming data into S3. This creates a data lake for further analysis.

Firehose can transform and analyze the data in flight before loading to S3 using Lambda. This enables real-time processing.

The data in S3 can be easily loaded into Amazon Redshift for interactive analysis at scale.

Kinesis auto scales to handle the high data volumes. Minimal effort is needed for infrastructure management.

upvoted 2 times

miki111 1 year, 10 months ago

Option D is the correct answer

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: D

A. This option utilizes S3 for data storage and EMR for analytics, Data Pipeline is not ideal service for real-time streaming data ingestion and processing. It is better suited for batch processing scenarios.

B. This option involves managing and scaling EC2, which adds operational overhead. It is also not real-time streaming solution. Additionally, use of Redshift for analyzing clickstream data might not be most efficient or cost-effective approach.

C. CloudFront is CDN service and is not designed for real-time data processing or analytics. While using Lambda to process data can be an option, it may not be most efficient solution for processing large volumes of clickstream data.

Therefore, collecting the data from Kinesis Data Streams, using Kinesis Data Firehose to transmit it to S3 data lake, and loading it into Redshift for analysis is the recommended approach. This combination provides scalable, real-time streaming solution with storage and analytics capabilities that can handle high volume of clickstream data.

upvoted 2 times

Rahulbit34 2 years, 1 month ago

Clickstream is the key - Answer is D

upvoted 1 times

PaoloRoma 2 years, 2 months ago

Selected Answer: A

I am going to be unpopular here and I'll go for A). Even if here are other services that offer a better experience, data Pipeline can do the job here. "you can use AWS Data Pipeline to archive your web server's logs to Amazon Simple Storage Service (Amazon S3) each day and then run a weekly Amazon EMR (Amazon EMR) cluster over those logs to generate traffic reports" <https://docs.aws.amazon.com/datapipeline/latest/DeveloperGuide/what-is-datapipeline.html> In the question there is no specific timing requirement for analytics. Also the EMR cluster job can be scheduled be executed daily.

Clickstream is the key - Answer is D

Option D) is a valid answer too, however with Amazon Redshift Streaming Ingestion "you can connect to Amazon Kinesis Data Streams data streams and pull data directly to Amazon Redshift without staging data in S3"

<https://aws.amazon.com/redshift/redshift-streaming-ingestion>. So in this scenario Kinesis Data Firehose and S3 are redundant.

upvoted 6 times

MutiverseAgent 1 year, 11 months ago

I think I agree with you, I does not make sense in option D) using Amazon Kinesis Data Firehose to transmit the data to an Amazon S3 data lake and then to Redshift, as you can send directly the data from Firehose to Redshift.

upvoted 2 times

juanrasus2 1 year, 8 months ago

Also the Kinesis family is related to real time or near real time services. This is not a requirement at all. We have to process data daily, but not need to do it in real time

upvoted 2 times

pentium75 1 year, 5 months ago

Question asks how to "transmit and process the clickstream data", NOT how to analyze it. This picture shows exactly scenario D:

Producer - Kinesis - Intermediate S3 bucket - Redshift

<https://d2908q01vomqb2.cloudfront.net/b6692ea5df920cad691c20319a6fffd7a4a766b8/2020/07/30/StreamTransformAnalyz-eKinesisLambdaRedshift1.png>

upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: D

Option D

upvoted 1 times

studis 2 years, 5 months ago

It is C.

The image in here <https://aws.amazon.com/kinesis/data-firehose/> shows how kinesis can send data collected to firehose who can send it to Redshift.

It is also possible to use an intermediary S3 bucket between firehose and redshift. See image in here

<https://aws.amazon.com/blogs/big-data/stream-transform-and-analyze-xml-data-in-real-time-with-amazon-kinesis-aws-lambda-and-amazon-redshift/>

upvoted 1 times

pentium75 1 year, 5 months ago

Makes sense, but this is D, not C

upvoted 1 times

sebasta 2 years, 6 months ago

Why not A?

You can collect data with AWS Data Pipeline and then analyze it with EMR. Whats wrong with this option?

upvoted 4 times

bearcandy 2 years, 6 months ago

It's not A, the wording is tricky! It says "to archive the data to S3" - there is no mention of archiving in the question, so it has to be D :)

upvoted 3 times

pentium75 1 year, 5 months ago

And, the the question is not asking about analyzing the data at all, just about "transmitting and processing".

upvoted 1 times

Wpcorgan 2 years, 6 months ago

D is correct

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #60

Topic 1


A company has a website hosted on AWS. The website is behind an Application Load Balancer (ALB) that is configured to handle HTTP and HTTPS separately. The company wants to forward all requests to the website so that the requests will use HTTPS.

What should a solutions architect do to meet this requirement?

- A. Update the ALB's network ACL to accept only HTTPS traffic.
- B. Create a rule that replaces the HTTP in the URL with HTTPS.
- C. Create a listener rule on the ALB to redirect HTTP traffic to HTTPS. **Most Voted**
- D. Replace the ALB with a Network Load Balancer configured to use Server Name Indication (SNI).

Correct Answer: C

Community vote distribution



Comments

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: C

C. Create a listener rule on the ALB to redirect HTTP traffic to HTTPS.

To meet the requirement of forwarding all requests to the website so that the requests will use HTTPS, a solutions architect can create a listener rule on the ALB that redirects HTTP traffic to HTTPS. This can be done by creating a rule with a condition that matches all HTTP traffic and a rule action that redirects the traffic to the HTTPS listener. The HTTPS listener should already be configured to accept HTTPS traffic and forward it to the target group.

upvoted 30 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option A. Updating the ALB's network ACL to accept only HTTPS traffic is not a valid solution because the network ACL is used to control inbound and outbound traffic at the subnet level, not at the listener level.

Option B. Creating a rule that replaces the HTTP in the URL with HTTPS is not a valid solution because this would not redirect the traffic to the HTTPS listener.

Option D. Replacing the ALB with a Network Load Balancer configured to use Server Name Indication (SNI) is not a valid solution because it would not address the requirement to redirect HTTP traffic to HTTPS.

upvoted 18 times

masetromain Highly Voted 2 years, 8 months ago

Selected Answer: C

Answer C :

https://docs.aws.amazon.com/fr_fr/elasticloadbalancing/latest/application/create-https-listener.html

<https://aws.amazon.com/fr/premiumsupport/knowledge-center/elb-redirect-http-to-https-using-alb/>

upvoted 15 times

satyaamm Most Recent 5 months, 1 week ago

Selected Answer: C

Redirecting the network on to HTTPS is the requirement here.

upvoted 1 times

cookieMr 8 months, 2 weeks ago

Selected Answer: C

A. Network ACLs operate at subnet level and control inbound and outbound traffic. Updating the network ACL alone will not enforce the redirection of HTTP to HTTPS.

B. This approach would require modifying application code or server configuration to perform URL rewrite. It is not an optimal solution as it adds complexity and potential maintenance overhead. Moreover, it does not leverage the ALB's capabilities for handling HTTP-to-HTTPS redirection.

D. While NLB can handle SSL/TLS termination using SNI for routing requests to different services, replacing the ALB solely to enforce HTTP-to-HTTPS redirection would be an unnecessary and more complex solution.

Therefore, the recommended approach is to create a listener rule on the ALB to redirect HTTP traffic to HTTPS. By configuring a listener rule, you can define a redirect action that automatically directs HTTP requests to their corresponding HTTPS versions.

upvoted 5 times

PaulGa 9 months ago

Selected Answer: C

Ans C - don't re-invent; just re-direct

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

<https://repost.aws/knowledge-center/elb-redirect-http-to-https-using-alb>

Steps 6-8 tells exactly how to do this:

"6. Select a load balancer, and then choose HTTP Listener.

7. Under Rules, choose View/edit rules.

8. Choose Edit Rule to modify the existing default rule to redirect all HTTP requests to HTTPS. Or, insert a rule between the existing rules (if appropriate for your use case)."

upvoted 3 times

Ruffyit 1 year, 7 months ago

C. Create a listener rule on the ALB to redirect HTTP traffic to HTTPS.

upvoted 3 times

AWSStudyBuddy 1 year, 7 months ago

Selected Answer: C

This solution meets all of the requirements:

Forward all requests to the website so that the requests will use HTTPS: The ALB can be configured to redirect all HTTP traffic to HTTPS.

The other options are not as good for this scenario:

A. Updating the ALB's network ACL to accept only HTTPS traffic will prevent users from accessing the website using HTTP.

B. Creating a rule that replaces the HTTP in the URL with HTTPS will not prevent users from accessing the website using HTTP.

D. Replacing the ALB with a Network Load Balancer configured to use Server Name Indication (SNI) is not necessary because the ALB can be configured to redirect all HTTP traffic to HTTPS.

upvoted 2 times

upvoted 3 times

Tom123456ac 1 year, 8 months ago

I hate this question description "The company wants to forward all requests to the website so that the requests will use HTTPS."

upvoted 3 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

The best solution is to create a listener rule on the Application Load Balancer (ALB) to redirect HTTP traffic to HTTPS (option C).

Here is why:

ALB listener rules allow you to redirect traffic from one listener port (e.g. 80 for HTTP) to another (e.g. 443 for HTTPS). This achieves the goal to forward all requests over HTTPS.

Network ACLs control traffic at the subnet level and cannot distinguish between HTTP and HTTPS requests to implement a redirect (option A incorrect).

Replacing HTTP with HTTPS in the URL happens at the client side. It does not redirect at the ALB (option B incorrect).

Network Load Balancers work at the TCP level and do not understand HTTP or HTTPS protocols. So they cannot redirect in this manner (option D incorrect).

upvoted 8 times

miki111 1 year, 10 months ago

Option C is the correct answer

upvoted 1 times

Abrar2022 2 years ago

A solutions architect should create listen rules to direct http traffic to https.

upvoted 1 times

cheese929 2 years, 1 month ago

Selected Answer: C

C is correct. Traffic redirection will solve it.

upvoted 2 times

elearningtakai 2 years, 2 months ago

Selected Answer: C

This rule can be created in the following way:

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. In the navigation pane, choose Load Balancers.
3. Select the ALB and choose Listeners.
4. Choose View/edit rules and then choose Add rule.
5. In the Add Rule dialog box, choose HTTPS.
6. In the Default action dialog box, choose Redirect to HTTPS.
7. Choose Save rules.

This listener rule will redirect all HTTP requests to HTTPS, ensuring that all traffic is encrypted.

upvoted 5 times

mell1222 2 years, 3 months ago

Selected Answer: C

Configure an HTTPS listener on the ALB: This step involves setting up an HTTPS listener on the ALB and configuring the security policy to use a secure SSL/TLS protocol and cipher suite.

Create a redirect rule on the ALB: The redirect rule should be configured to redirect all incoming HTTP requests to HTTPS. This can be done by creating a redirect rule that redirects HTTP requests on port 80 to HTTPS requests on port 443.

Update the DNS record: The DNS record for the website should be updated to point to the ALB's DNS name, so that all traffic is routed through the ALB.

Verify the configuration: Once the configuration is complete, the website should be tested to ensure that all requests are being redirected to HTTPS. This can be done by accessing the website using HTTP and verifying that the request is redirected to HTTPS.

upvoted 2 times

career360guru 2 years, 5 months ago

Selected Answer: C

Option C

upvoted 1 times

Shasha1 2 years, 6 months ago

C

To redirect HTTP traffic to HTTPS, a solutions architect should create a listener rule on the ALB to redirect HTTP traffic to HTTPS. Option A is not correct because network ACLs do not have the ability to redirect traffic. Option B is not correct because it does not redirect traffic, it only replaces the URL. Option D is not correct because a Network Load Balancer does not have the ability to handle HTTPS traffic.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #61

Topic 1

A company is developing a two-tier web application on AWS. The company's developers have deployed the application on an Amazon EC2 instance that connects directly to a backend Amazon RDS database. The company must not hardcode database credentials in the application. The company must also implement a solution to automatically rotate the database credentials on a regular basis.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Store the database credentials in the instance metadata. Use Amazon EventBridge (Amazon CloudWatch Events) rules to run a scheduled AWS Lambda function that updates the RDS credentials and instance metadata at the same time.
- B. Store the database credentials in a configuration file in an encrypted Amazon S3 bucket. Use Amazon EventBridge (Amazon CloudWatch Events) rules to run a scheduled AWS Lambda function that updates the RDS credentials and the credentials in the configuration file at the same time. Use S3 Versioning to ensure the ability to fall back to previous values.
- C. Store the database credentials as a secret in AWS Secrets Manager. Turn on automatic rotation for the secret. Attach the required permission to the EC2 role to grant access to the secret. **Most Voted**
- D. Store the database credentials as encrypted parameters in AWS Systems Manager Parameter Store. Turn on automatic rotation for the encrypted parameters. Attach the required permission to the EC2 role to grant access to the encrypted parameters.

Correct Answer: C

Community vote distribution

C (100%)

Comments

KVK16 **Highly Voted** 2 years, 8 months ago

Selected Answer: C

Secrets manager supports Autorotation unlike Parameter store.
upvoted 25 times

JesseeS 2 years, 7 months ago

Parameter store does not support autorotation.
upvoted 10 times

updated 16 times

Buruguduystunstugudunstuy Highly Voted 2 years, 5 months ago

Selected Answer: C

The correct solution is C. Store the database credentials as a secret in AWS Secrets Manager. Turn on automatic rotation for the secret. Attach the required permission to the EC2 role to grant access to the secret.

AWS Secrets Manager is a service that enables you to easily rotate, manage, and retrieve database credentials, API keys, and other secrets throughout their lifecycle. By storing the database credentials as a secret in Secrets Manager, you can ensure that they are not hardcoded in the application and that they are automatically rotated on a regular basis. To grant the EC2 instance access to the secret, you can attach the required permission to the EC2 role. This will allow the application to retrieve the secret from Secrets Manager as needed.

upvoted 16 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option A, storing the database credentials in the instance metadata and using a Lambda function to update them, would not meet the requirement of not hardcoding the credentials in the application.

Option B, storing the database credentials in an encrypted S3 bucket and using a Lambda function to update them, would also not meet this requirement, as the application would still need to access the credentials from the configuration file.

Option D, storing the database credentials as encrypted parameters in AWS Systems Manager Parameter Store, would also not meet this requirement, as the application would still need to access the encrypted parameters in order to use them.

upvoted 10 times

DatNX Most Recent 3 months, 3 weeks ago

Selected Answer: C

- automatically rotate credentials on regular basis -> AWS Secret Manager

upvoted 2 times

Kyleskii 4 months, 2 weeks ago

Selected Answer: C

Keyword: database credentials.

AWS secrets managers will handle it.

upvoted 1 times

satyaamm 5 months, 1 week ago

Selected Answer: C

Rotating the credentials is a feature provided by AWS Secrets Manager

upvoted 1 times

PaulGa 9 months ago

Selected Answer: C

Ans C - Secrets Manager, provides rotation - and also a lot more API calls

upvoted 1 times

soufiyane 1 year, 2 months ago

Selected Answer: C

parameter store does not have auto rotation

upvoted 1 times

Atul6969 1 year, 3 months ago

Selected Answer: C

test kjlshfjkh jfskjfnkj kj bskjfb kj kjs bfkjs b kjf

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

Secrets Manager is purpose built for this scenario

secrets manager is purpose built for this scenario

AB are wrong and insecure way of doing this

D Parameter store with encrypted string can be used for this but is not ideal choice and AFAIK it does not support automatic rotation without extra programming

upvoted 1 times

1Alpha1 1 year, 5 months ago

Selected Answer: C

C - "Auto Rotation"

upvoted 1 times

Ruffyt 1 year, 7 months ago

AWS Secrets Manager is a service that enables you to easily rotate, manage, and retrieve database credentials, API keys, and other secrets throughout their lifecycle. By storing the database credentials as a secret in Secrets Manager, you can ensure that they are not hardcoded in the application and that they are automatically rotated on a regular basis. To grant the EC2 instance access to the secret, you can attach the required permission to the EC2 role.

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

Storing the credentials in AWS Secrets Manager and enabling automatic rotation meets the requirements with the least operational overhead. The EC2 instance role just needs permission to access the secret, and Secrets Manager handles rotating the credentials automatically on a schedule.

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: C

Key Autorotation = AWS Secrets Manager

upvoted 2 times

miki111 1 year, 10 months ago

Option C is the right answer.

upvoted 1 times

cookieMr 1 year, 11 months ago

Selected Answer: C

Storing the credentials in Secrets Manager provides dedicated and secure management. With automatic rotation enabled, Secrets Manager handles the credential updates automatically. Attaching the necessary permissions to the EC2 role allows the application to securely access the secret.

This approach minimizes operational overhead and provides a secure and managed solution for credential management.

upvoted 2 times

Bmarodi 2 years ago

Selected Answer: C

The solution that meets the requirements with the least operational overhead, is option C.

upvoted 1 times

Bmarodi 2 years ago

Selected Answer: C

My choice is c.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #62

Topic 1


A company is deploying a new public web application to AWS. The application will run behind an Application Load Balancer (ALB). The application needs to be encrypted at the edge with an SSL/TLS certificate that is issued by an external certificate authority (CA). The certificate must be rotated each year before the certificate expires.

What should a solutions architect do to meet these requirements?

- A. Use AWS Certificate Manager (ACM) to issue an SSL/TLS certificate. Apply the certificate to the ALB. Use the managed renewal feature to automatically rotate the certificate.
- B. Use AWS Certificate Manager (ACM) to issue an SSL/TLS certificate. Import the key material from the certificate. Apply the certificate to the ALB. Use the managed renewal feature to automatically rotate the certificate.
- C. Use AWS Certificate Manager (ACM) Private Certificate Authority to issue an SSL/TLS certificate from the root CA. Apply the certificate to the ALB. Use the managed renewal feature to automatically rotate the certificate.
- D. Use AWS Certificate Manager (ACM) to import an SSL/TLS certificate. Apply the certificate to the ALB. Use Amazon EventBridge (Amazon CloudWatch Events) to send a notification when the certificate is nearing expiration. Rotate the certificate manually. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

Sinaneos **Highly Voted** 2 years, 8 months ago

Selected Answer: D

It's a third-party certificate, hence AWS cannot manage renewal automatically. The closest thing you can do is to send a notification to renew the 3rd party certificate.

upvoted 61 times

mabotega **Highly Voted** 2 years, 7 months ago

Selected Answer: D

It is D, because ACM does not manage the renewal process for imported certificates. You are responsible for monitoring the expiration date of your imported certificates and for renewing them before they expire.

Check this question on the link below:

Q: What types of certificates can I create and manage with ACM?

https://www.amazonaws.cn/en/certificate-manager/faqs/#Managed_renewal_and_deployment

upvoted 22 times

Dharmarajan Most Recent 4 months, 2 weeks ago

Selected Answer: D

D is the most appropriate. C seems like it might work, but the word "Private" makes it ineligible answer. Although I would say is to have ACM renew the public certificate automatically per the renewal method supported: https://www.amazonaws.cn/en/certificate-manager/faqs/#Managed_renewal_and_deployment

Thanks @mabotega for the link to the doc.

upvoted 1 times

satyaamm 5 months, 1 week ago

Selected Answer: D

Since the certificate is external it needs to be imported and thus D is the most suitable option.

upvoted 1 times

MrAnhQuan 5 months, 3 weeks ago

Selected Answer: D

It's D, because ACM doesn't manage the renewal process for imported certificates.

upvoted 2 times

Tjazz04 6 months ago

Selected Answer: D

Ans. D

Keyword: issued by an external certificate authority

ACM cannot rotate certificate automatically if from external CA.

upvoted 3 times

PaulGa 9 months ago

Selected Answer: D

Ans D - hint: we're importing the certificate "...Use AWS Certificate Manager (ACM) to import an SSL/TLS certificate"

upvoted 3 times

Anji195 1 year, 4 months ago

Yes it's D. Here is a clear explanation.

Imported certificates – If you want to use a third-party certificate with Amazon CloudFront, Elastic Load Balancing, or Amazon API Gateway, you may import it into ACM using the AWS Management Console, AWS CLI, or ACM APIs. ACM can not renew imported certificates, but it can help you manage the renewal process. You are responsible for monitoring the expiration date of your imported certificates and for renewing them before they expire. You can use ACM CloudWatch metrics to monitor the expiration dates of an imported certificates and import a new third-party certificate to replace an expiring one.

https://www.amazonaws.cn/en/certificate-manager/faqs/#Managed_renewal_and_deployment

upvoted 5 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

"certificate that is issued by an external certificate authority (CA)"

AB will create a new certificate in AWS

C will also create a new certificate but this is not what PCA are for * =

(<https://docs.aws.amazon.com/privateca/latest/userguide/PcaWelcome.html>)

D: Import the certificate is correct answer

upvoted 1 times

1Alpha1 1 year, 5 months ago

Selected Answer: D

D - "External CA" --> 'Update Manually'

upvoted 1 times

xdkonorek2 1 year, 7 months ago

Selected Answer: A

internal CA are typically trusted only within the organization unless you manually distribute and trust the root certificate elsewhere

external CA:

Certificates from a well-known external CA are trusted by most browsers and systems by default

<https://docs.aws.amazon.com/acm/latest/userguide/acm-certificate.html>

"Public certificates that you request through ACM are obtained from Amazon Trust Services, an Amazon managed public certificate authority (CA). ... Any browser, application, or OS that includes the Amazon or Starfield roots will trust public certificates obtained from ACM."

The answer is A, different story if they said external certificate

upvoted 3 times

Ruffyt 1 year, 7 months ago

: What types of certificates can I create and manage with ACM?

https://www.amazonaws.cn/en/certificate-manager/faqs/#Managed_renewal_and_deployment

upvoted 1 times

est3la21 1 year, 9 months ago

answer is D

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: D

The key points are:

Obtain certificate from external CA, not ACM

Import the external certificate into ACM

Apply imported certificate to the ALB

Set up EventBridge rule to trigger notification on certificate expiration

Manually renew and rotate the external certificate each year.

upvoted 2 times

miki111 1 year, 10 months ago

Option D is the right answer.

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: D

D: With this approach, you import the third-party certificate into ACM, which allows you to centrally manage and apply it to the ALB. By configuring CloudWatch Events, you can receive notifications when the certificate is close to expiring, prompting you to manually initiate the rotation process.

A & B: These options assume that the SSL/TLS certificate can be issued directly by ACM. However, since the requirement specifies that the certificate should be issued by an external certificate authority (CA), this option is not suitable.

C: ACM Private Certificate Authority is used when you want to create your own private CA and issue certificates from it. It does not support certificates issued by external CAs. Therefore, this option is not suitable for the given requirement.

upvoted 3 times

Router 1 year, 12 months ago

D is correct, since it's an external certificate

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #63

Topic 1

A company runs its infrastructure on AWS and has a registered base of 700,000 users for its document management application. The company intends to create a product that converts large .pdf files to .jpg image files. The .pdf files average 5 MB in size. The company needs to store the original files and the converted files. A solutions architect must design a scalable solution to accommodate demand that will grow rapidly over time.

Which solution meets these requirements MOST cost-effectively?

A. Save the .pdf files to Amazon S3. Configure an S3 PUT event to invoke an AWS Lambda function to convert the files to .jpg format and store them back in Amazon S3. **Most Voted**


B. Save the .pdf files to Amazon DynamoDB. Use the DynamoDB Streams feature to invoke an AWS Lambda function to convert the files to .jpg format and store them back in DynamoDB.

C. Upload the .pdf files to an AWS Elastic Beanstalk application that includes Amazon EC2 instances, Amazon Elastic Block Store (Amazon EBS) storage, and an Auto Scaling group. Use a program in the EC2 instances to convert the files to .jpg format. Save the .pdf files and the .jpg files in the EBS store.

D. Upload the .pdf files to an AWS Elastic Beanstalk application that includes Amazon EC2 instances, Amazon Elastic File System (Amazon EFS) storage, and an Auto Scaling group. Use a program in the EC2 instances to convert the file to .jpg format. Save the .pdf files and the .jpg files in the EBS store.

Correct Answer: A

Community vote distribution



Comments

ArielSchivo **Highly Voted** 2 years, 7 months ago

Selected Answer: A

Option A. Elastic BeanStalk is expensive, and DocumentDB has a 400KB max to upload files. So Lambda and S3 should be the one.

upvoted 52 times

raffaello44 2 years, 7 months ago

is lambda scalable as an EC2 ?

upvoted 5 times

EtherealBagel 1 year, 6 months ago

lambda has near infinite scale
upvoted 6 times

rob74 2 years, 7 months ago

In addition to this Lambda is paid only when used....
upvoted 7 times

mrbottomwood 2 years, 5 months ago

I'm thinking when you wrote DocumentDB you meant it as DynamoDB...yes?
upvoted 6 times

benjl 2 years, 5 months ago

Yes, DynamoDB has 400KB limit for the item.
<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/ServiceQuotas.html>
upvoted 8 times

cookieMr Highly Voted 1 year, 11 months ago

Selected Answer: A

B. Using DynamoDB for storing and processing large .pdf files would not be cost-effective due to storage and throughput costs associated with DynamoDB.

C. Using Elastic Beanstalk with EC2 and EBS storage can work, but it may not be most cost-effective solution. It involves managing the underlying infrastructure and scaling manually.

D. Similar to C, using Elastic Beanstalk with EC2 and EFS storage can work, but it may not be most cost-effective solution. EFS is a shared file storage service and may not provide optimal performance for conversion process, especially as demand and file sizes increase.

A. leverages Lambda and the scalable and cost-effective storage of S3. With Lambda, you only pay for actual compute time used during the file conversion, and S3 provides durable and scalable storage for both .pdf files and .jpg files. The S3 PUT event triggers Lambda to perform conversion, eliminating need to manage infrastructure and scaling, making it most cost-effective solution for this scenario.

upvoted 10 times

PaulGa Most Recent 9 months ago

Selected Answer: A

Ans A - a simple get and PUT back to S3 bucket. At "...average 5 MB" the returned .jpeg files should be smaller than standard S3
upvoted 1 times

shasi07 11 months, 3 weeks ago

BeanStack is expensive solution, and dynamoDB have a limitation of 400KB max to upload files. So Lambda an S3 should be the one.

upvoted 1 times

Uzbekistan 1 year, 2 months ago

Given the company's requirement for access to both AWS and on-premises file storage with minimum latency and no significant changes to existing file access patterns, the most suitable option is:

A. Deploy and configure Amazon FSx for Windows File Server on AWS. Move the on-premises file data to FSx for Windows File Server. Reconfigure the workloads to use FSx for Windows File Server on AWS.

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

S3 is the only scalable option for such a large user base in cost effective way.
BCD can work but will be extremely costly

upvoted 1 times

Ruffyt 1 year, 7 months ago

B. Using DynamoDB for storing and processing large .pdf files would not be cost-effective due to storage and throughput costs associated with DynamoDB.

C. Using Elastic Beanstalk with EC2 and EBS storage can work, but it may not be most cost-effective solution. It involves managing the underlying infrastructure and scaling manual

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: A

Option A is the most cost-effective solution that meets the requirements. Here is why:

Storing the PDFs in Amazon S3 is inexpensive and scalable storage.

Using S3 events to trigger Lambda functions to do the file conversion is a serverless approach that scales automatically. No need to manage EC2 instances.

Lambda usage is charged only for compute time used, which is cost-efficient for spiky workloads like this.

Storing the converted JPGs back in S3 keeps the storage scalable and cost-effective.

upvoted 4 times

RDX19 1 year, 10 months ago

Selected Answer: A

Option A is right answer since Dynamo DB has size limitations.

upvoted 2 times

miki111 1 year, 10 months ago

Option A is the right answer.

upvoted 1 times

Bmarodi 2 years ago

Selected Answer: A

The solution meets these requirements most cost-effectively is option A.

upvoted 1 times

Bmarodi 2 years ago

Selected Answer: A

I think the best solution is A.

Ref. <https://s3.amazonaws.com/doc/s3-developer-guide/RESTObjectPUT.html>

upvoted 1 times

Abrar2022 2 years ago

Since this requires a cost-effect solution then you can use Lambda to convert pdf files to jpeg and store them on S3. Lambda is serverless, so only pay when you use it and automatically scales to cope with demand.

upvoted 1 times

srirajav 2 years, 1 month ago

if Option A is correct, however storing the data back to the same S3, wont it cause infinite looping, it's not best practice right storing a object that is processed by Lambda function to the same S3 bucket, it has chances to cause infinite Loop and then if the option B cant we increase the limits of Dynamo DB requesting AWS?

upvoted 3 times

bedwal2020 2 years, 1 month ago

In question, it is never mentioned that the jpg files will also be stored in same s3 bucket. We can have different s3 buckets right ?

upvoted 2 times

cheese929 2 years, 1 month ago

Selected Answer: A

Answer A is the most cost effective solution that meets the requirement

upvoted 1 times

channn 2 years, 2 months ago

Selected Answer: A

Key words: MOST cost-effectively, so S3 + Lambda
upvoted 1 times

SilentMilli 2 years, 5 months ago

Selected Answer: A

This solution will meet the company's requirements in a cost-effective manner because it uses a serverless architecture with AWS Lambda to convert the files and store them in S3. The Lambda function will automatically scale to meet the demand for file conversions and S3 will automatically scale to store the original and converted files as needed.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #64

Topic 1

A company has more than 5 TB of file data on Windows file servers that run on premises. Users and applications interact with the data each day.


The company is moving its Windows workloads to AWS. As the company continues this process, the company requires access to AWS and on-premises file storage with minimum latency. The company needs a solution that minimizes operational overhead and requires no significant changes to the existing file access patterns. The company uses an AWS Site-to-Site VPN connection for connectivity to AWS.

What should a solutions architect do to meet these requirements?

- A. Deploy and configure Amazon FSx for Windows File Server on AWS. Move the on-premises file data to FSx for Windows File Server. Reconfigure the workloads to use FSx for Windows File Server on AWS.
- B. Deploy and configure an Amazon S3 File Gateway on premises. Move the on-premises file data to the S3 File Gateway. Reconfigure the on-premises workloads and the cloud workloads to use the S3 File Gateway.
- C. Deploy and configure an Amazon S3 File Gateway on premises. Move the on-premises file data to Amazon S3. Reconfigure the workloads to use either Amazon S3 directly or the S3 File Gateway, depending on each workload's location.
- D. Deploy and configure Amazon FSx for Windows File Server on AWS. Deploy and configure an Amazon FSx File Gateway on premises. Move the on-premises file data to the FSx File Gateway. Configure the cloud workloads to use FSx for Windows File Server on AWS. Configure the on-premises workloads to use the FSx File Gateway. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

sba21 **Highly Voted** 2 years, 8 months ago

Selected Answer: D

<https://www.examtopycs.com/discussions/amazon/view/83281-exam-aws-certified-solutions-architect-associate-saa-c02/>
upvoted 26 times

MutiverseAgent 1 year, 11 months ago

Agree answer is D)

Requirements are:

- "Users and applications interact with the data each day"
- "the company requires access to AWS and on-premises file storage with minimum latency"

Explanation: Answer A) will work with the same on-prem <> aws latency as in answer D) as both use the VPN Connection. Having said this, by using an Amazon FSx File Gateway on premise as the D) scenario mentioned, all users will have a great benefit on using the cache that the FSx File Gateway has on their daily workloads. And that is part of the requirements: "users", "each day", "latency"

upvoted 5 times

MrAWS Highly Voted 2 years, 4 months ago

D IS WRONG - Its used for caching. you cannot 'Move the on-premises file data to the FSx File Gateway.' which is stated in answer D. It pretty sure AWS employee's are spamming this site with the wrong answers intentionally.

upvoted 20 times

Rabbit117 1 year ago

I think D is correct. The FSx File gateway will look like any other Windows file share to the on-prem clients so when they write, or "Copy" data to the share it will be be cached locally and also written to the FSx service in the cloud.

<https://docs.aws.amazon.com/filegateway/latest/filefsxw/file-gateway-fsx-concepts.html>

upvoted 3 times

DarthVaper 1 year, 8 months ago

What's the problem with it being a cache? They did say "the company requires access to AWS and on-premises file storage with minimum latency."

Not discarding what you said but what's wrong here?

upvoted 4 times

dsshahu01 1 year, 3 months ago

The problem with cache is - It needs to refreshed, which is an overkill for a migration project

The cache refresh requires another solution since the users/applications interact with it daily which means it does modify often.

upvoted 2 times

chagantik90 1 year, 6 months ago

you don't move data to gateway, its cached when people use those files from Fsx server

upvoted 1 times

senrix Most Recent 3 weeks, 5 days ago

Selected Answer: D

D will address the "latency" requirement

upvoted 1 times

Dharmarajan 4 months, 2 weeks ago

Selected Answer: D

On reconsidering, A might be incorrect. First part of te answer looks ok, but later it says "Configure apps to use FSx in AWS". So that doesnt make sense - what about the on prem data?

So perhaps its better to simply move all data to FSx on AWS. I think none of the choices are perfect. So D comes close.

upvoted 1 times

Dharmarajan 4 months, 2 weeks ago

Selected Answer: A

D is incorrect choice.The documentation describes the FSx File Gateway as:

"FSx for Windows File Server makes all the shares on the server available as shares on your Amazon FSx File Gateway." - so it is simply a fronting for the FSx Windows File Server, not an independent storage FSx system.

So I would choose A as a more appropriate response.

upvoted 1 times

dariar 4 months, 2 weeks ago

Selected Answer: A

You can't migrate data to FSx File Gateway

Migrate your files to FSx for Windows File Server. For more information, see Migrating existing file storage to FSx for Windows File Server.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: A

Ans A - it can only be A or D. With D introducing unnecessary operations for an eventual AWS migration

upvoted 1 times

iyiola_daniel 9 months, 1 week ago

Selected Answer: A

The question says the company is moving it's workload from windows to AWS. Plus D involves operational overhead.

upvoted 3 times

MatAlves 10 months ago

Selected Answer: D

Bad question.

A - Doesn't meet the "low latency" requirement due to VPN connection.

D - It's not possible to "move file data to FSx File Gateway". Instead, it provides improved latency with local caching.

B and C - wrong, since the question requires "no significant changes to the existing file access patterns"

I'd go with D, but not 100% happy with it.

upvoted 4 times

hieunt.hus 11 months ago

Selected Answer: A

A. Deploy and configure Amazon FSx for Windows File Server on AWS. Move the on-premises file data to FSx for Windows File Server. Reconfigure the workloads to use FSx for Windows File Server on AWS.

upvoted 2 times

jatric 11 months, 2 weeks ago

Selected Answer: D

keyword as company continue to migrate workload from om-prem to cloud

upvoted 2 times

Abhi2024 1 year, 1 month ago

I would go with option A as question mentions "solution that minimizes operational overhead" . Adding FSx File Gateway would add more complexity. They already have Site to Site VPN. This way min change required.

upvoted 2 times

ManikRoy 1 year, 1 month ago

Selected Answer: D

Will go with option D as option A is missing the Fsx File gateway component deployed on-premise which will provide the low latency access.

upvoted 3 times

3680113 1 year ago

Option D is Correct, D provides low latency options for both of premium and the cloud loads.

upvoted 2 times

3680113 1 year ago

on-prem

upvoted 1 times

stormbreaker7 1 year, 1 month ago

Selected Answer: A

min. operational overhead

min. operational overhead.

upvoted 1 times

Uzbekistan 1 year, 2 months ago

Selected Answer: A

Option D suggests deploying both Amazon FSx for Windows File Server on AWS and an Amazon FSx File Gateway on premises. While this option may provide a solution for accessing on-premises file data in AWS, it introduces additional complexity and potential overhead that may not be necessary given the company's existing AWS Site-to-Site VPN connectivity.

Here are some reasons why Option D may not be the most suitable choice:

Complexity, Cost, Redundancy and Operational Overhead.

upvoted 1 times

3680113 1 year ago

Both Amazon FSx file gateway on-prem and Amazon FSx for windows file server goes hand in hand. So there is no way you can use the file gateway on-Prem and not use the Fsx for windows file server. Your conclusion of adding more overhead does not hold. Hence D is the perfect answer.

Configure the cloud workloads to use FSx for Windows File Server on AWS. Configure the on-premises workloads to use the FSx File Gateway.

upvoted 1 times

Rabbit117 1 year ago

A VPN or DX connection is required for FSx File gateway. Option D provides low maintenance as FSx for Windows is a managed service and the file gateway will provide the low latency access required by the on-prem services.

upvoted 1 times

LIORAGE 1 year, 2 months ago

Selected Answer: D

D: FSX File gateway is necessary for communication between on-premise and AWS FSX for window

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

Windows File Servers + Preserve compatibility so BC is wrong due to S3

A does not provide on-premise access and suggest to move the files which is wrong as company wants to keep on-prem access

D meets all the requirements.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #65

Topic 1


A hospital recently deployed a RESTful API with Amazon API Gateway and AWS Lambda. The hospital uses API Gateway and Lambda to upload reports that are in PDF format and JPEG format. The hospital needs to modify the Lambda code to identify protected health information (PHI) in the reports.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use existing Python libraries to extract the text from the reports and to identify the PHI from the extracted text.
- B. Use Amazon Textract to extract the text from the reports. Use Amazon SageMaker to identify the PHI from the extracted text.
- C. Use Amazon Textract to extract the text from the reports. Use Amazon Comprehend Medical to identify the PHI from the extracted text. **Most Voted**
- D. Use Amazon Rekognition to extract the text from the reports. Use Amazon Comprehend Medical to identify the PHI from the extracted text.

Correct Answer: C

Community vote distribution



Comments

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: C

The correct solution is C: Use Amazon Textract to extract the text from the reports. Use Amazon Comprehend Medical to identify the PHI from the extracted text.

Option C: Using Amazon Textract to extract the text from the reports, and Amazon Comprehend Medical to identify the PHI from the extracted text, would be the most efficient solution as it would involve the least operational overhead. Textract is specifically designed for extracting text from documents, and Comprehend Medical is a fully managed service that can accurately identify PHI in medical text. This solution would require minimal maintenance and would not incur any additional costs beyond the usage fees for Textract and Comprehend Medical.

upvoted 21 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option A: Using existing Python libraries to extract the text and identify the PHI from the text would require the hospital to maintain and update the libraries as needed. This would involve operational overhead in terms of keeping the libraries up to date and

debugging any issues that may arise.

Option B: Using Amazon SageMaker to identify the PHI from the extracted text would involve additional operational overhead in terms of setting up and maintaining a SageMaker model, as well as potentially incurring additional costs for using SageMaker.

Option D: Using Amazon Rekognition to extract the text from the reports would not be an effective solution, as Rekognition is primarily designed for image recognition and would not be able to accurately extract text from PDF or JPEG files.

upvoted 8 times

satyaamm Most Recent 5 months, 1 week ago

Selected Answer: C

Textract is more suitable than Rekognition as it is build to scan text documents and Comprehend helps identify PHI.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: C

Ans C - Textract to 'read' data; Comprehend to assess whether its PHI

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

Textract = Extract text from PDF/images

Comprehend Medical = PHI

ABD are wrong products for this requirement so won't achieve the results

upvoted 3 times

djgodzilla 1 year, 5 months ago

Selected Answer: C

Both Rekognition and Textract possess the ability to detect text within images, yet they are optimized for differing applications.

Rekognition specializes in identifying text located spatially within an image, for instance, words displayed on street signs, t-shirts, or license plates. Its typical use cases encompass visual search, content filtering, deriving insights from content, among others. However, it's not the ideal choice for images containing more than 100 words, as this exceeds its limitation.

On the other hand, Textract is tailored more towards processing documents and PDFs, offering a comprehensive suite for Optical Character Recognition (OCR). It proves useful in scenarios involving financial reports, medical records, receipts, ID documents, and more.

upvoted 4 times

Ruffyit 1 year, 7 months ago

The correct solution is C: Use Amazon Textract to extract the text from the reports. Use Amazon Comprehend Medical to identify the PHI from the extracted text.

Option C: Using Amazon Textract to extract the text from the reports, and Amazon Comprehend Medical to identify the PHI from the extracted text, would be the most efficient solution as it would involve the least operational overhead. Textract is specifically designed for extracting text from documents, and Comprehend Medical is a fully managed service that can accurately identify PHI in medical text. This solution would require minimal maintenance and would not incur any additional costs beyond the usage fees for Textract and Comprehend Medical.

upvoted 2 times

AWSSStudyBuddy 1 year, 7 months ago

Selected Answer: C

• Amazon Textract: This program is made to extract text and data from scanned documents, such as pictures and PDFs. It helps to retain the formatting of the report by automatically extracting text while preserving the document's layout.

Identifying and extracting medical information, including protected health information (PHI), from unstructured text is the specialty of Amazon Comprehend Medical. Medical entities that are frequently included in reporting on healthcare, such as ailments, drugs, and more, can be recognized by it.

upvoted 2 times

Chiquitabandita 1 year, 9 months ago

with the choices here. I would go with C, but if offered, I would use amazon textract for the text and use Macie to do the scanning

of text files, not comprehend.

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

Here's why:

Amazon Textract has built-in support to extract text from PDFs and images, eliminating the need to build this yourself with Python libraries.

Amazon Comprehend Medical has pre-trained machine learning models to identify PHI entities out-of-the-box, avoiding the need to train your own SageMaker model.

Using these fully managed AWS services minimizes operational overhead of maintaining machine learning models yourself.

upvoted 1 times

miki111 1 year, 10 months ago

Option C is the right answer.

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: C

C leverages capabilities of Textract, which is a service that automatically extracts text and data from documents, including PDF and JPEG. By using Textract, hospital can extract text content from reports without need for additional custom code or libraries.

Once text is extracted, hospital can then use Comprehend Medical, a natural language processing service specifically designed for medical text, to analyze and identify PHI. It can recognize medical entities such as medical conditions, treatments, and patient information.

A. suggests using existing Python libraries, which would require hospital to develop and maintain custom code for text extraction and PHI identification.

B and D involve using Textract along with SageMaker or Rekognition, respectively, for PHI identification. While these options could work, they introduce additional complexity by incorporating machine learning models and training.

upvoted 2 times

channn 2 years, 2 months ago

Key word: hospital!

upvoted 1 times

alexiscloud 2 years, 2 months ago

Answer C:

upvoted 1 times

Chirantan 2 years, 5 months ago

Selected Answer: C

Amazon Textract is a machine learning (ML) service that automatically extracts text, handwriting, and data from scanned documents.

upvoted 3 times

career360guru 2 years, 5 months ago

Selected Answer: C

Option C

upvoted 1 times

SONA_M 2 years, 6 months ago

WHY OPTION D IS WRONG

upvoted 1 times

s_fun 2 years, 5 months ago

D is wrong only because Amazon Rekognition doesn't read text, only explicit image contents.

upvoted 3 times

mj61 2 years, 4 months ago

B/C you use TextTract to extract text not Rekognition.
upvoted 1 times

k1kavi1 2 years, 6 months ago

Selected Answer: C

Agreed
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #66

Topic 1

A company has an application that generates a large number of files, each approximately 5 MB in size. The files are stored in Amazon S3. Company policy requires the files to be stored for 4 years before they can be deleted. Immediate accessibility is always required as the files contain critical business data that is not easy to reproduce. The files are frequently accessed in the first 30 days of the object creation but are rarely accessed after the first 30 days.

Which storage solution is MOST cost-effective?

- A. Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 Glacier 30 days from object creation. Delete the files 4 years after object creation.
- B. Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA) 30 days from object creation. Delete the files 4 years after object creation.
- C. Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) 30 days from object creation. Delete the files 4 years after object creation. **Most Voted**
- D. Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) 30 days from object creation. Move the files to S3 Glacier 4 years after object creation.

Correct Answer: C*Community vote distribution***Comments****Six_Fingered_Jose** **Highly Voted** 2 years, 7 months ago**Selected Answer: C**

i think C should be the answer here,

> Immediate accessibility is always required as the files contain critical business data that is not easy to reproduce

If they do not explicitly mention that they are using Glacier Instant Retrieval, we should assume that Glacier -> takes more time to retrieve and may not meet the requirements

upvoted 92 times

JayBee65 2 years, 6 months ago

You can make that assumption, but I think it would be wrong to make it. It does not state they are not using Glacier Instant Retrieval, and it's use would be the logical choice in this question, so I'm going for A

upvoted 8 times

syh_rapha 2 years, 6 months ago

I think his assumption is correct because if you go to AWS documentation (<https://aws.amazon.com/s3/storage-classes/glacier/>) they clearly mention: "S3 Glacier Flexible Retrieval (formerly S3 Glacier)". So since this question doesn't specify the S3 Glacier class, then it would default to flexible retrieval (which ofc is not equal to Instant Retrieval).

upvoted 11 times

slackbot 1 year, 9 months ago

why everybody assumed files must be deleted after 4 years. they said files "can" be deleted, and not "must" be deleted. ideally store the files in glacier after 4 years

upvoted 3 times

wearrexdzw3123 1 year, 7 months ago

Because it requires the lowest cost

upvoted 4 times

sandordini 1 year, 2 months ago

"With S3 Glacier Instant Retrieval, you can save up to 68% on storage costs compared to using the S3 Standard-Infrequent Access (S3 Standard-IA) storage class, when your data is accessed once per quarter." Repeat: once per quarter! Plus if you keep all the data forever, even if on Glacier, it will cost more over time, than storing on a higher-cost storage for a period of time (here: 4 yrs). A does not delete.

upvoted 3 times

sandordini 1 year, 2 months ago

Sorry! A also deletes... My bad... Still. once a quarter.

upvoted 1 times

Kumaran1508 2 years ago

Yeah, Correct answer is C

Because even if you assume the glacier class as Instant Retrieval. As per the Instant Retrieval class the immediate availability is only once per quarter. But in question it is clearly mentioned that the files should be immediately available anytime.

upvoted 6 times

ninjawrz Highly Voted 2 years, 8 months ago

Selected Answer: A

Most COST EFFECTIVE

A: S3 Glacier Instant Retrieval is a new storage class that delivers the fastest access to archive storage, with the same low latency and high-throughput performance as the S3 Standard and S3 Standard-IA storage classes. You can save up to 68 percent on storage costs as compared with using the S3 Standard-IA storage class when you use the S3 Glacier Instant Retrieval storage class and pay a low price to retrieve data.

upvoted 31 times

babayomi 9 months ago

COST EFFECTIVE IS THE DIAL WORD HERE. Answer is A

upvoted 2 times

PoolDead 10 months, 2 weeks ago

Glacier has high retrieval costs and long retrieval times, which does not meet the requirement for immediate accessibility.

upvoted 2 times

Pamban 2 years, 7 months ago

"Immediate accessibility is always required as the files contain critical business data that is not easy to reproduce" is the key sentence. answer is C.

upvoted 7 times

JayBee65 2 years, 6 months ago

But S3 Glacier Instant Retrieval "is designed for rarely accessed data that still needs immediate access in performance."

but S3 Glacier Instant Retrieval "is designed for rarely accessed data that still needs immediate access in performance-sensitive use cases", so it offers lower cost and instant retrieval, so A

upvoted 2 times

Bala75krish 2 years, 4 months ago

I agree with your key sentence..but the one zone infrequent doesn't fit for critical business and it is used for recreate..

upvoted 1 times

wearrexdzw3123 1 year, 7 months ago

Glaciers usually take some time to retrieve

upvoted 1 times

bora4motion Most Recent 1 month, 1 week ago

Selected Answer: C

"Immediate accessibility is always required " and then "Cost effective"- that leaves us with C.

upvoted 1 times

bora4motion 1 month, 2 weeks ago

Selected Answer: C

I'm going with C - the files might be required instantly even if rarely accessed.

upvoted 1 times

KeshavP 3 months, 1 week ago

Selected Answer: C

S3 Glacier Instant Retrieval is more expensive than S3 Standard-IA for both storage and retrieval requests. S3 Standard-IA is 68% cheaper for storage than S3 Standard, while S3 Glacier Instant Retrieval is only 46% cheaper

upvoted 1 times

Vandaman 3 months, 3 weeks ago

Selected Answer: A

Answer is A. Files are rarely accessed after 30 days and they require the MOST cost effective solution. Then the answer is Glacier.

upvoted 1 times

kyd0nix 4 months, 1 week ago

Selected Answer: C

If Glacier class is not specified, I could assume it is referring to S3 Glacier Flexible Retrieval (formerly S3 Glacier). Then, the retrieval of data is not immediate as requested. The best option would be glacier Instant Retrieval, but it's not specified... so the best answer that fits is C.

upvoted 1 times

babyalrtabs 4 months, 2 weeks ago

Selected Answer: C

C is the answer. While both "S3 Standard-Infrequent Access" and "S3 One Zone-Infrequent Access" are storage classes on AWS S3 designed for data accessed infrequently, the key difference is that "S3 One Zone-IA" stores data in a single Availability Zone, making it cheaper but with a higher risk of data loss if that zone fails, whereas "S3 Standard-IA" spreads data across multiple Availability Zones for increased redundancy and availability; making it a better choice for critical data that needs to be readily accessible even if rarely accessed.

upvoted 1 times

satyaammm 5 months, 1 week ago

Selected Answer: C

S3 IA is the most suitable here.

upvoted 1 times

PSH123 6 months, 2 weeks ago

Selected Answer: B

onezone IA is cheap more than standard IA

upvoted 4 times

Ode7d1b 6 months, 3 weeks ago

Option C is the most cost-effective solution because it provides immediate access for the first 30 days and optimizes for cost by transitioning the files to S3 Standard-IA after that period. It also complies with the company's retention policy.

upvoted 1 times

PaPi_z 8 months, 3 weeks ago

Selected Answer: C

Immediate accessibility

upvoted 3 times

anminfang 9 months ago

It should be C. The critical point is Immediate accessibility, only S3-IA can provide this.

upvoted 2 times

MatAlves 10 months ago

Selected Answer: C

"Immediate accessibility is always required". It can't be a, since Glacier Instant Retrieval has a limit of times the data can be accessed.

"S3 Glacier Retrieval is the lowest cost archive storage with milliseconds retrieval for rarely accessed data. It is ideal for data that is accessed once or twice per quarter, and that requires immediate access."

upvoted 2 times

MatAlves 10 months ago

If you do so more than ~4 times per year, you would pay more for Glacier Instant Retrieval than for S3 Standard.

upvoted 1 times

jatric 11 months, 2 weeks ago

Selected Answer: C

immediately accessible for 4 years but rarely use - S3 standard IA.

Correct answer is C

upvoted 2 times

MoSaleh8890 4 months, 3 weeks ago

Why not A, with Glacier Instant Retrieval?

upvoted 1 times

Duckydoo 1 year ago

Selected Answer: C

Keep in mind that "Glacier Flexible Retrieval" was previously named just "Glacier" (before they added Instant Retrieval). Regardless, they should have made it clear which Glacier class they are referring to since which one you select/assume can impact the correctness.

upvoted 3 times

lofzee 1 year ago

Selected Answer: B

I am going to go with B. Reasoning =

A = even if Glacier Instant Retrieval was used here, you can only use that quick retrieval once a quarter which isn't what the prompt is asking for.

B - One Zone is the cheapest option, and the prompt does not ask for high availability.

C / D - see reasoning for B

upvoted 5 times

Duckydoo 1 year ago

It does mention that the data is not easily reproducible, which implies HA, so that eliminates 1Z in this case.

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #67

Topic 1


A company hosts an application on multiple Amazon EC2 instances. The application processes messages from an Amazon SQS queue, writes to an Amazon RDS table, and deletes the message from the queue. Occasional duplicate records are found in the RDS table. The SQS queue does not contain any duplicate messages.

What should a solutions architect do to ensure messages are being processed once only?

- A. Use the CreateQueue API call to create a new queue.
- B. Use the AddPermission API call to add appropriate permissions.
- C. Use the ReceiveMessage API call to set an appropriate wait time.
- D. Use the ChangeMessageVisibility API call to increase the visibility timeout. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

KVK16 **Highly Voted** 2 years, 8 months ago

Selected Answer: D

In case of SQS - multi-consumers if one consumer has already picked the message and is processing, in meantime other consumer can pick it up and process the message there by two copies are added at the end. To avoid this the message is made invisible from the time its picked and deleted after processing. This visibility timeout is increased according to max time taken to process the message

upvoted 53 times

JayBee65 2 years, 6 months ago

To add to this "The VisibilityTimeout in SQS is a time frame that the message can be hidden so that no others can consume it except the first consumer who calls the ReceiveMessageAPI." The API ChangeMessageVisibility changes this value.

upvoted 21 times

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: D

To ensure that messages are being processed only once, a solutions architect should use the ChangeMessageVisibility API call to increase the visibility timeout which is Option D.

The visibility timeout determines the amount of time that a message received from an SQS queue is hidden from other consumers while the message is being processed. If the processing of a message takes longer than the visibility timeout, the message will become visible to other consumers and may be processed again. By increasing the visibility timeout, the solutions architect can ensure that the message is not made visible to other consumers until the processing is complete and the message can be safely deleted from the queue.

Option A (Use the CreateQueue API call to create a new queue) would not address the issue of duplicate message processing.

Option B (Use the AddPermission API call to add appropriate permissions) is not relevant to this issue.

Option C (Use the ReceiveMessage API call to set an appropriate wait time) is also not relevant to this issue.

upvoted 18 times

karbob 2 years, 5 months ago

not relevant to this issue. ??? what is added value

upvoted 3 times

Buruguduystunstugudunstuy 2 years, 3 months ago

Option B (Use the AddPermission API call to add appropriate permissions) is not relevant to this issue because it deals with setting permissions for accessing an SQS queue, which is not related to preventing duplicate records in the RDS table.

Option C (Use the ReceiveMessage API call to set an appropriate wait time) is not relevant to this issue because it is related to configuring how long the ReceiveMessage API call should wait for new messages to arrive in the SQS queue before returning an empty response. It does not address the issue of duplicate records in the RDS table.

upvoted 8 times

Dharmarajan Most Recent 4 months, 2 weeks ago

Selected Answer: D

I chose D as well, I understand the rationale behind it. But I think it is named incorrectly. Should this not be "InvisibilityTimeout" instead of "VisibilityTimeout"?

If the message becomes visible again after a timeout, because SQS assumes it failed processing - then it remains invisible and after timeout it becomes visible again... So that's where my point of contention is.

upvoted 2 times

Rcosmos 5 months ago

Selected Answer: D

Explicação:

Causa provável do problema:

O problema ocorre porque o tempo de visibilidade (visibility timeout) configurado na fila do Amazon SQS é muito curto. Quando uma instância EC2 lê uma mensagem da fila e começa a processá-la, se o processamento não for concluído antes que o tempo de visibilidade expire, a mensagem volta para a fila e pode ser processada novamente por outra instância. Isso leva a registros duplicados na tabela RDS.

upvoted 1 times

satyaamm 5 months, 1 week ago

Selected Answer: D

Visibility timeout is used to manage multiple users.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: D

Ans D - once the message is selected in the queue 'hide' it from other users/apps so it is only processed once, using API 'ChangeMessageVisibility'

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

AB: Irrelevant

C: This is for long polling not for execution

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-short-and-long-polling.html#sqs-long-polling>

long polling

D: Visibility is correct fix issue because over here other SQS clients are seeing the same message back in the queue when the previous processor is taking longer than expected to process the message

upvoted 1 times

Subhrangsu 1 year, 8 months ago

I also opt for D, but asking does increasing MessageVisibilityTimeout good always?

upvoted 2 times

awsgeek75 1 year, 5 months ago

This parameter is the timeout after which the SQS assumes that the processing has failed and makes the item visible to other queue processors. It is normal to increase the timeout if it takes longer to process items.

upvoted 1 times

miki111 1 year, 10 months ago

Option D is the right answer.

upvoted 1 times

cookieMr 1 year, 11 months ago

Selected Answer: D

The visibility timeout is the duration during which SQS prevents other consumers from receiving and processing the same message. By increasing the visibility timeout, you allow more time for the processing of a message to complete before it becomes visible to other consumers.

Option A, creating a new queue, does not address the issue of concurrent processing and duplicate records. It would only create a new queue, which is not necessary for solving the problem.

Option B, adding permissions, also does not directly address the issue of duplicate records. Permissions are necessary for accessing the SQS queue but not for preventing concurrent processing.

Option C, setting an appropriate wait time using the ReceiveMessage API call, does not specifically prevent duplicate records. It can help manage the rate at which messages are received from the queue but does not address the issue of concurrent processing.

upvoted 4 times

cheese929 2 years, 1 month ago

Selected Answer: D

D is correct

upvoted 1 times

alexiscloud 2 years, 2 months ago

Answer D:

visibility timeout beings when amazon SQS return a message

upvoted 1 times

test_devops_aws 2 years, 2 months ago

Selected Answer: D

D = ChangeMessageVisibility

upvoted 1 times

dev1978 2 years, 4 months ago

In theory, between reception and changing visibility, you can have multiple consumers. Question is not good as it won't guarantee not executing twice.

upvoted 1 times

techhb 2 years, 5 months ago

Selected Answer: D

Increaseing visibility timeout makes sure message is not visible for time taken to process the message.

upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: D

Option D

upvoted 1 times

Wpcorgan 2 years, 6 months ago

D is correct

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #68

Topic 1

A solutions architect is designing a new hybrid architecture to extend a company's on-premises infrastructure to AWS. The company requires a highly available connection with consistent low latency to an AWS Region. The company needs to minimize costs and is willing to accept slower traffic if the primary connection fails.

What should the solutions architect do to meet these requirements?

- A. Provision an AWS Direct Connect connection to a Region. Provision a VPN connection as a backup if the primary Direct Connect connection fails. **Most Voted**
- B. Provision a VPN tunnel connection to a Region for private connectivity. Provision a second VPN tunnel for private connectivity and as a backup if the primary VPN connection fails.
- C. Provision an AWS Direct Connect connection to a Region. Provision a second Direct Connect connection to the same Region as a backup if the primary Direct Connect connection fails.
- D. Provision an AWS Direct Connect connection to a Region. Use the Direct Connect failover attribute from the AWS CLI to automatically create a backup connection if the primary Direct Connect connection fails.

Correct Answer: A

Community vote distribution

A (93%)

Other (7%)

Comments

KVK16 **Highly Voted** 2 years, 8 months ago

Selected Answer: A

Direct Connect + VPN best of both
upvoted 21 times

mabotega **Highly Voted** 2 years, 7 months ago

Selected Answer: A

Direct Connect goes through 1 Gbps, 10 Gbps or 100 Gbps and the VPN goes up to 1.25 Gbps.

<https://docs.aws.amazon.com/whitepapers/latest/aws-vpc-connectivity-options/aws-direct-connect-vpn.html>
upvoted 19 times

akeed **Most Recent** 2 months, 4 weeks ago

Selected Answer: D

Highly available means the system is up and running most of the time, not 100% of the time. Therefore, minor outages are acceptable.

upvoted 1 times

Dharmarajan 4 months, 2 weeks ago

Selected Answer: A

Since the direct connect takes a lot of time to setup and is much more expensive, the point of consideration is - if those are not in the criteria, we assume Direct connection is alright.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: A

And A - as primary use Direct Connect to provide a solid connection. Company is not too worried about backup (other than it works) so use cheaper VPN Site-to-Site.

upvoted 2 times

48cd959 1 year, 2 months ago

Selected Answer: A

Ans A-

Direct Connect because company needs consistent connection.

As a back up, company wants cheaper solution so VPN Site to Site connection should be okay.

upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

HA low latency + minimize cost + acceptable slow traffic if primary fails

B: VPN tunnel will be slow

C: 2 direct connect will be expensive

D: Backup connection for what?

A: Direct connect + VPN as a backup works

upvoted 5 times

Ruffyt 1 year, 7 months ago

A highly available connection with consistent low latency = AWS Direct Connect

Minimize costs and accept slower traffic if the primary connection fails = VPN connection

upvoted 6 times

benacert 1 year, 9 months ago

A is the right choice to save cost

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: A

Highly available connectivity using Direct Connect for consistent low latency and high throughput.

Cost optimization by using a VPN as a slower, lower cost backup for when Direct Connect fails.

Automatic failover to the VPN when Direct Connect fails.

upvoted 3 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: A

A highly available connection with consistent low latency = AWS Direct Connect

Minimize costs and accept slower traffic if the primary connection fails = VPN connection

upvoted 1 times

hsinchang 1 year, 10 months ago

Selected Answer: A

Slower traffic when primary fails, so the backup plan needs a cheaper solution, and the primary requires high performance, so

Δ

upvoted 1 times

oguzbeliren 1 year, 10 months ago

Even though, there are a lots of variable affecting the cost of the connection, VPN connection is cheaper than the Direct Connect most of the time since VPN Connection doesn't require any dedicated physical circuit involved.

upvoted 1 times

miki111 1 year, 10 months ago

Option A is the right answer.

upvoted 1 times

cookieMr 1 year, 11 months ago

Selected Answer: A

Options B and C propose using multiple VPN connections for private connectivity and as backups. While VPNs can serve as backups, they may not provide the same level of consistent low latency and high availability as Direct Connect connections. Additionally, provisioning multiple VPN tunnels can increase operational complexity and costs.

Option D suggests using the Direct Connect failover attribute from the AWS CLI to automatically create a backup connection if the primary Direct Connect connection fails. While this approach can be automated, it does not provide the same level of immediate failover capabilities as having a separate backup connection in place.

Therefore, option A, provisioning an AWS Direct Connect connection to a Region and provisioning a VPN connection as a backup, is the most suitable solution that meets the company's requirements for connectivity, cost-effectiveness, and high availability.

upvoted 4 times

th3k33n 2 years, 1 month ago

Selected Answer: A

highly available - > direct connect because connection can go up to 10GBPs and VPN 1.5GBPs as backup

upvoted 2 times

linux_admin 2 years, 2 months ago

Selected Answer: A

Option A is the correct solution to meet the requirements of the company. Provisioning an AWS Direct Connect connection to a Region will provide a private and dedicated connection with consistent low latency. As the company requires a highly available connection, a VPN connection can be provisioned as a backup if the primary Direct Connect connection fails. This approach will minimize costs and provide the required level of availability.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #69

Topic 1

A company is running a business-critical web application on Amazon EC2 instances behind an Application Load Balancer. The EC2 instances are in an Auto Scaling group. The application uses an Amazon Aurora PostgreSQL database that is deployed in a single Availability Zone. The company wants the application to be highly available with minimum downtime and minimum loss of data.

Which solution will meet these requirements with the LEAST operational effort?

A. Place the EC2 instances in different AWS Regions. Use Amazon Route 53 health checks to redirect traffic. Use Aurora PostgreSQL Cross-Region Replication.


B. Configure the Auto Scaling group to use multiple Availability Zones. Configure the database as Multi-AZ. Configure an Amazon RDS Proxy instance for the database. **Most Voted**

C. Configure the Auto Scaling group to use one Availability Zone. Generate hourly snapshots of the database. Recover the database from the snapshots in the event of a failure.

D. Configure the Auto Scaling group to use multiple AWS Regions. Write the data from the application to Amazon S3. Use S3 Event Notifications to launch an AWS Lambda function to write the data to the database.

Correct Answer: B

Community vote distribution



Comments

SilentMilli **Highly Voted** 2 years, 5 months ago

Selected Answer: B

By configuring the Auto Scaling group to use multiple Availability Zones, the application will be able to continue running even if one Availability Zone goes down. Configuring the database as Multi-AZ will also ensure that the database remains available in the event of a failure in one Availability Zone. Using an Amazon RDS Proxy instance for the database will allow the application to automatically route traffic to healthy database instances, further increasing the availability of the application. This solution will meet the requirements for high availability with minimal operational effort.

upvoted 37 times

KVK16 **Highly Voted** 2 years, 8 months ago

Selected Answer: B

RDS Proxy for Aurora <https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html>

upvoted 8 times

AshishDhole Most Recent 5 months ago

Selected Answer: B

Highly Available = Multi AZ

DR = Multiple Region

upvoted 2 times

Rcosmos 5 months ago

Selected Answer: B

Conclusão:

A configuração B oferece alta disponibilidade tanto para as instâncias do EC2 quanto para o banco de dados Aurora PostgreSQL com o mínimo esforço operacional. Isso atende aos requisitos de tempo de inatividade mínimo e perda mínima de dados.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: B

Ans B - good explanation by SilentMilli (1yr 8mth ago). The database Proxy is key.

upvoted 2 times

SaurabhTiwari1 9 months, 3 weeks ago

Selected Answer: B

Keyword - Higly available

Use multiple availability zones

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

A: Different region doesn't help

C: Would have made sense if it wasn't restricting to one AZ.

D: Regions + S3 + Lambda = Operational effort extreme

B: Although not entirely sure how RDS Proxy helps because it is for connection pooling but it is the only workable solution using multi AZ

upvoted 3 times

dkw2342 1 year, 3 months ago

One of the benefits of Amazon RDS Proxy is that it can improve application recovery time after database failovers. While RDS Proxy supports both MySQL as well as PostgreSQL engines, in this post, we will use a MySQL test workload to demonstrate how RDS Proxy reduces client recovery time after failover by up to 79% for Amazon Aurora MySQL and by up to 32% for Amazon RDS for MySQL.

-> contributes to minimum downtime req

upvoted 1 times

dkw2342 1 year, 3 months ago

PS: <https://aws.amazon.com/blogs/database/improving-application-availability-with-amazon-rds-proxy/>

upvoted 1 times

MiniYang 1 year, 6 months ago

Selected Answer: A

The company wants to minimize costs and is willing to accept slower traffic if the primary connection fails, it may be tempted to choose a VPN connection as a backup, in which case the answer is A. Cost-Effectiveness: VPN connections are generally more economical than AWS Direct Connect, especially for low to moderate bandwidth needs.

Backup connection: A VPN connection can serve as a more cost-effective backup if the primary Direct Connect connection fails, even if it may be slower. Acceptance of slower traffic: The question clearly states that the company is willing to accept slower traffic if the primary connection fails, which implies a tolerance for connection speeds.

upvoted 1 times

pentium75 1 year, 5 months ago

This is about the previous question ;)

upvoted 6 times

asulhi 1 year, 8 months ago

Selected Answer: B

ASG and MultiAZ is the best answer

upvoted 1 times

benacert 1 year, 9 months ago

B is the right answer

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

Option B requires the least operational effort to meet the high availability and minimum downtime/data loss requirements.

The key points are:

Use an Auto Scaling group across multiple AZs for high availability of the EC2 instances.

Configure the Aurora DB as Multi-AZ for high availability, automatic failover, and minimum data loss.

Use RDS Proxy for connection pooling to the DB for performance

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: B

Highly available, Minimum downtime and Minimum loss of data = Auto Scaling group on Multi-AZ, Database on Multi-AZ, Amazon RDS Proxy.

upvoted 1 times

miki111 1 year, 10 months ago

Option B is the right answer.

upvoted 1 times

hiepdz98 1 year, 11 months ago

Selected Answer: B

B is correct answer

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: B

A. This approach provides geographic redundancy, it introduces additional complexity and operational effort, including managing replication, handling latency, and potentially higher data transfer costs.

C. While snapshots can be used for data backup and recovery, they do not provide real-time failover capabilities and can result in significant data loss if a failure occurs between snapshots.

D. While this approach offers some decoupling and scalability benefits, it adds complexity to the data flow and introduces additional overhead for data processing.

In comparison, option B provides a simpler and more streamlined solution by utilizing multiple AZs, Multi-AZ configuration for the database, and RDS Proxy for improved connection management. It ensures high availability, minimal downtime, and minimum loss of data with the least operational effort.

upvoted 6 times

Abrar2022 2 years ago

@Wajif the reason why it's not A is because the question mentions High availability and nothing to do with region. You can achieve HA without spanning multiple regions. Also B is incorrect because ALB are region specific and span across multiple AZ with that specific region (not cross region)

upvoted 1 times

UnluckyDucky 2 years, 4 months ago

OmuckYDuckY 2 years, 7 months ago

Selected Answer: B

RDS Proxy is fully managed by AWS for RDS/Aurora. It is auto-scaling and highly available by default.
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #70


Topic 1

A company's HTTP application is behind a Network Load Balancer (NLB). The NLB's target group is configured to use an Amazon EC2 Auto Scaling group with multiple EC2 instances that run the web service.

The company notices that the NLB is not detecting HTTP errors for the application. These errors require a manual restart of the EC2 instances that run the web service. The company needs to improve the application's availability without writing custom scripts or code.

What should a solutions architect do to meet these requirements?

- A. Enable HTTP health checks on the NLB, supplying the URL of the company's application.
- B. Add a cron job to the EC2 instances to check the local application's logs once each minute. If HTTP errors are detected, the application will restart.
- C. Replace the NLB with an Application Load Balancer. Enable HTTP health checks by supplying the URL of the company's application. Configure an Auto Scaling action to replace unhealthy instances. **Most Voted**
- D. Create an Amazon Cloud Watch alarm that monitors the UnhealthyHostCount metric for the NLB. Configure an Auto Scaling action to replace unhealthy instances when the alarm is in the ALARM state.

Correct Answer: C*Community vote distribution***Comments****123jh10** **Highly Voted** 2 years, 7 months ago**Selected Answer: C**

I would choose A, as NLB supports HTTP and HTTPS Health Checks, BUT you can't put any URL (as proposed), only the node IP addresses.

So, the solution is C.

upvoted 32 times

jerry1 2 months, 3 weeks ago

I think even http health check is feasible, option A does not provide any suggestion when health check fails

upvoted 1 times

Ack3rman 2 years, 6 months ago

can you elaborate more pls
upvoted 3 times

BlueVolcano1 2 years, 4 months ago

NLBs support HTTP, HTTPS and TCP health checks:
<https://docs.aws.amazon.com/elasticloadbalancing/latest/network/target-group-health-checks.html> (check HealthCheckProtocol)

But NLBs only accept either selecting EC2 instances or IP addresses directly as targets. You can't provide a URL to your endpoints, only a health check path (if you're using HTTP or HTTPS health checks).

upvoted 11 times

km142646 2 years, 1 month ago

What's the difference between endpoint URL and health check path?

upvoted 2 times

majubmo 2 years ago

A URL includes the hostname. The health check path is only the path portion. For example,
URL = <https://i-0123456789abcdef.us-west-2.compute.internal/index.html>
health check path= /index.html

upvoted 15 times

ArielSchivo Highly Voted 2 years, 7 months ago

Selected Answer: C

Option C. NLB works at Layer 4 so it does not support HTTP/HTTPS. The replacement for the ALB is the best choice.

upvoted 20 times

BlueVolcano1 2 years, 4 months ago

That's incorrect. NLB does support HTTP and HTTPS (and TCP) health checks.
<https://docs.aws.amazon.com/elasticloadbalancing/latest/network/target-group-health-checks.html>

There just isn't an answer option that reflects that. My guess is that the question and/or answer options are outdated.

upvoted 7 times

Agatsume 3 weeks ago

NLB doesnt support HTTP health checks..only TCP

upvoted 1 times

AwsAbhiKumar Most Recent 4 months, 2 weeks ago

Selected Answer: C

NLB only supports TCP-level health checks, it cannot detect HTTP errors

upvoted 2 times

satyaamm 5 months, 1 week ago

Selected Answer: C

ALB is more suitable for HTTP requests.

upvoted 2 times

PaulGa 9 months ago

Selected Answer: C

Ans C - Yup, something's happening at the Application level so replace NLB with ALB

upvoted 4 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

NLB is for network errors and low level traffic stuff

ALB is for application so C is the only realistic option here

upvoted 8 times

kel2023 1 year, 5 months ago

Selected Answer: A

NLB does support HTTP/HTTPS Health Checks.

I saw other people comments, it seems like the question were rephrased. The comments were highlighting "application URL", but I don't see words on the question.

upvoted 3 times

ignajtpolandstrong 1 year, 5 months ago

Selected Answer: C

You can use HTTP/HTTPS ONLY when Target is ALB.

By default it is TCP.

<https://docs.aws.amazon.com/elasticloadbalancing/latest/network/target-group-health-checks.html#health-check-settings>
HealthCheckProtocol

The protocol the load balancer uses when performing health checks on targets. The possible protocols are HTTP, HTTPS, and TCP. The default is the TCP protocol. If the target type is ALB, the supported health check protocols are HTTP and HTTPS.

upvoted 3 times

tom_cruise 1 year, 8 months ago

Selected Answer: C

ALB allows you to specify the path which helps to check the error. NLB cannot do that.

upvoted 4 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

The key points are:

Use an Application Load Balancer (ALB) instead of a Network Load Balancer (NLB) since ALBs support HTTP health checks.

Configure HTTP health checks on the ALB to monitor the application health.

Use an Auto Scaling action triggered by the ALB health checks to automatically replace unhealthy instances.

upvoted 4 times

miki111 1 year, 10 months ago

Option C is the right answer.

upvoted 1 times

cookieMr 1 year, 11 months ago

Selected Answer: C

A. NLB, but NLB's health checks are designed for TCP/UDP protocols and lack the advanced features specific to HTTP applications provided by ALB.

B. This approach involves custom scripting and manual intervention, which contradicts the requirement of not writing custom scripts or code.

D. Since the NLB does not detect HTTP errors, relying solely on the UnhealthyHostCount metric may not accurately capture the health of the application instances.

Therefore, C is the recommended choice for improving the application's availability without custom scripting or code. By replacing the NLB with an ALB, enabling HTTP health checks, and configuring Auto Scaling to replace unhealthy instances, the company can ensure that only healthy instances are serving traffic, enhancing the application's availability automatically.

upvoted 10 times

Gizmo2022 6 months, 3 weeks ago

Thank you for explaining this, I'm currently studying for the exam and your explanation was helpful.

upvoted 1 times

Abrar2022 2 years ago

Replace the NLB (layer 4 udp and tcp) with an Application Load Balancer - ALB (layer 7) supports http and https requests.

upvoted 1 times

datz 2 years, 2 months ago

Selected Answer: C

must be C

Application availability: NLB cannot assure the availability of the application. This is because it bases its decisions solely on network and TCP-layer variables and has no awareness of the application at all. Generally, NLB determines availability based on the ability of a server to respond to ICMP ping or to correctly complete the three-way TCP handshake. ALB goes much deeper and is capable of determining availability based on not only a successful HTTP GET of a particular page but also the verification that the content is as was expected based on the input parameters.

upvoted 1 times

datz 2 years, 2 months ago

Also A doesn't offer what bellow in C offers...

Configure an Auto Scaling action to replace unhealthy instances

upvoted 1 times

Tony1980 2 years, 4 months ago

Answer is C

A solution architect can use Amazon EC2 Auto Scaling health checks to automatically detect and replace unhealthy instances in the EC2 Auto Scaling group. The health checks can be configured to check the HTTP errors returned by the application and terminate the unhealthy instances. This will ensure that the application's availability is improved, without requiring custom scripts or code.

upvoted 1 times

aakashkumar1999 2 years, 4 months ago

I will go with A as Network load balancer supports HTTP and HTTPS health checks, maybe the answer is outdated.

upvoted 2 times

pentium75 1 year, 5 months ago

But you'd need to check the health of the individual nodes, NOT "the URL of the company's application" which points to the Load Balancer.

upvoted 2 times

John_Zhuang 2 years, 5 months ago

Selected Answer: C

<https://medium.com/awesome-cloud/aws-difference-between-application-load-balancer-and-network-load-balancer-cb8b6cd296a4>

As NLB does not support HTTP health checks, you can only use ALB to do so.

upvoted 1 times

BlueVolcano1 2 years, 4 months ago

That's incorrect. NLB does support HTTP and HTTPS (and TCP) health checks.

<https://docs.aws.amazon.com/elasticloadbalancing/latest/network/target-group-health-checks.html>

Just a general tip: Medium is not a reliable resource. Anyone can create content there. Rely only on official AWS documentation.

upvoted 4 times

pentium75 1 year, 5 months ago

But you'd need to check the health of the individual nodes, NOT "the URL of the company's application" which points to the Load Balancer (as mentioned in A).

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #71

Topic 1

A company runs a shopping application that uses Amazon DynamoDB to store customer information. In case of data corruption, a solutions architect needs to design a solution that meets a recovery point objective (RPO) of 15 minutes and a recovery time objective (RTO) of 1 hour.

What should the solutions architect recommend to meet these requirements?

- A. Configure DynamoDB global tables. For RPO recovery, point the application to a different AWS Region.
- B. Configure DynamoDB point-in-time recovery. For RPO recovery, restore to the desired point in time. **Most Voted**
- C. Export the DynamoDB data to Amazon S3 Glacier on a daily basis. For RPO recovery, import the data from S3 Glacier to DynamoDB.
- D. Schedule Amazon Elastic Block Store (Amazon EBS) snapshots for the DynamoDB table every 15 minutes. For RPO recovery, restore the DynamoDB table by using the EBS snapshot.

Correct Answer: B

Community vote distribution

B (100%)

Comments

123jh10 **Highly Voted** 2 years, 7 months ago

Selected Answer: B

A - DynamoDB global tables provides multi-Region, and multi-active database, but it not valid "in case of data corruption". In this case, you need a backup. This solutions isn't valid.

****B**** - Point in Time Recovery is designed as a continuous backup juts to recover it fast. It covers perfectly the RPO, and probably the RTO. <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/PointInTimeRecovery.html>

C - A daily export will not cover the RPO of 15min.

D - DynamoDB is serverless... so what are these EBS snapshots taken from???

upvoted 49 times

LionelSid 2 years, 4 months ago

Yes, it is possible to take EBS snapshots of a DynamoDB table. The process for doing this involves the following steps:

Create a new Amazon Elastic Block Store (EBS) volume from the DynamoDB table.

Stop the DynamoDB service on the instance.

Detach the EBS volume from the instance.

Create a snapshot of the EBS volume.

Reattach the EBS volume to the instance.

Start the DynamoDB service on the instance.

You can also use AWS Data pipeline to automate the above process and schedule regular snapshots of your DynamoDB table.

Note that, if your table is large and you want to take a snapshot of it, it could take a long time and consume a lot of bandwidth, so it's recommended to use the Global Tables feature from DynamoDB in order to have a Multi-region and Multi-master DynamoDB table, and you can snapshot each region separately.

upvoted 4 times

piavik 2 years, 2 months ago

What is "DynamoDB service on the instance" ?

upvoted 2 times

pentium75 1 year, 5 months ago

DynamoDB is a native cloud service, there is no "instance" that you could "stop", or detach an "EBS volume" from.

upvoted 1 times

Buruguduystunstugudunstuy Highly Voted 2 years, 5 months ago

Selected Answer: B

The best solution to meet the RPO and RTO requirements would be to use DynamoDB point-in-time recovery (PITR). This feature allows you to restore your DynamoDB table to any point in time within the last 35 days, with a granularity of seconds. To recover data within a 15-minute RPO, you would simply restore the table to the desired point in time within the last 35 days.

To meet the RTO requirement of 1 hour, you can use the DynamoDB console, AWS CLI, or the AWS SDKs to enable PITR on your table. Once enabled, PITR continuously captures point-in-time copies of your table data in an S3 bucket. You can then use these point-in-time copies to restore your table to any point in time within the retention period.

CORRECT

Option B. Configure DynamoDB point-in-time recovery. For RPO recovery, restore to the desired point in time.

upvoted 14 times

Buruguduystunstugudunstuy 2 years, 5 months ago

WRONG

Option A (configuring DynamoDB global tables) would not meet the RPO requirement, as global tables are designed to replicate data to multiple regions for high availability, but they do not provide a way to restore data to a specific point in time.

Option C (exporting data to S3 Glacier) would not meet the RPO or RTO requirements, as S3 Glacier is a cold storage service with a retrieval time of several hours.

Option D (scheduling EBS snapshots) would not meet the RPO requirement, as EBS snapshots are taken on a schedule, rather than continuously. Additionally, restoring a DynamoDB table from an EBS snapshot can take longer than 1 hour, so it would not meet the RTO requirement.

upvoted 7 times

satyaamm Most Recent 5 months, 1 week ago

Selected Answer: B

DynamoDB Point-In-Time-Recovery is the most suitable option here for recovering an object.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: B

Ans B - just use the built-in DynamoDB PITR... simples...

upvoted 1 times

ManikRoy 1 year, 1 month ago

Selected Answer: B

With point-in-time recovery, you can restore that table to any point in time during the last 35 days. After you enable point-in-time recovery, you can restore to any point in time from five minutes before the current time until 35 days ago. DynamoDB maintains incremental backups of your table.

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

A: Scalability across regions which is not required

C: Glacier exports and backup restore won't meet 1 hour RPO time

D EBS for DynamoDB table? Sounds impractical

B: DynamoDB point-in-time recovery is for this scenario.

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/PointInTimeRecovery.html>

upvoted 2 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

The best option to meet the RPO of 15 minutes and RTO of 1 hour is B) Configure DynamoDB point-in-time recovery. For RPO recovery, restore to the desired point in time.

The key points:

DynamoDB point-in-time recovery can restore to any point in time within the last 35 days. This supports an RPO of 15 minutes. Restoring from a point-in-time backup meets the 1 hour RTO.

Point-in-time recovery is specifically designed to restore DynamoDB tables with second-level granularity.

upvoted 1 times

cookieMr 1 year, 11 months ago

A. Global tables provide multi-region replication for disaster recovery purposes, they may not meet the desired RPO of 15 minutes without additional configuration and potential data loss.

C. Exporting and importing data on a daily basis does not align with the desired RPO of 15 minutes.

D. EBS snapshots can be used for data backup, they are not directly applicable to DynamoDB and cannot provide the desired RPO and RTO without custom implementation.

In comparison, option B utilizing DynamoDB's built-in point-in-time recovery functionality provides the most straightforward and effective solution for meeting the specified RPO of 15 minutes and RTO of 1 hour. By enabling PITR and restoring the table to the desired point in time, the company can recover the customer information with minimal data loss and within the required time frame.

upvoted 3 times

Abrar2022 2 years ago

The answer is in the question. Read the question again!!! Option B. Configure DynamoDB point-in-time recovery. For RPO recovery, restore to the desired point in time.

upvoted 1 times

[Removed] 2 years, 1 month ago

If there is anyone who is willing to share his/her contributor access, then please write to vinaychethi99@gmail.com

upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: B

Option B

upvoted 1 times

Shasha1 2 years, 6 months ago

B is correct

DynamoDB point-in-time recovery allows the solutions architect to recover the DynamoDB table to a specific point in time, which would meet the RPO of 15 minutes. This feature also provides an RTO of 1 hour, which is the desired recovery time objective for the application. Additionally, configuring DynamoDB point-in-time recovery does not require any additional infrastructure or operational effort, making it the best solution for this scenario.

Option D is not correct because scheduling Amazon EBS snapshots for the DynamoDB table every 15 minutes would not meet the RPO or RTO requirements. While EBS snapshots can be used to recover data from a DynamoDB table, they are not designed to provide real-time data protection or recovery capabilities

upvoted 1 times

Wpcorgan 2 years, 6 months ago

B is correct

upvoted 1 times

SimonPark 2 years, 7 months ago

Selected Answer: B

B is the answer

upvoted 1 times

BoboChow 2 years, 7 months ago

Selected Answer: B

I think DynamoDB global tables also work here, but Point in Time Recovery is a better choice

upvoted 1 times

Kikiokiki 2 years, 7 months ago

I THINK B.

<https://dynobase.dev/dynamodb-point-in-time-recovery/>

upvoted 1 times

priya2224 2 years, 8 months ago

answer is D

upvoted 1 times

123jh10 2 years, 7 months ago

DynamoDB is serverless, so no storage snapshots available. <https://aws.amazon.com/dynamodb/>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #72

Topic 1

A company runs a photo processing application that needs to frequently upload and download pictures from Amazon S3 buckets that are located in the same AWS Region. A solutions architect has noticed an increased cost in data transfer fees and needs to implement a solution to reduce these costs.

How can the solutions architect meet this requirement?

- A. Deploy Amazon API Gateway into a public subnet and adjust the route table to route S3 calls through it.
- B. Deploy a NAT gateway into a public subnet and attach an endpoint policy that allows access to the S3 buckets.
- C. Deploy the application into a public subnet and allow it to route through an internet gateway to access the S3 buckets.
- D. Deploy an S3 VPC gateway endpoint into the VPC and attach an endpoint policy that allows access to the S3 buckets.

Most Voted

Correct Answer: D

Community vote distribution

D (100%)

Comments

Buruguduystunstugudunstuy Highly Voted 2 years, 5 months ago

Selected Answer: D

CORRECT

The correct answer is Option D. Deploy an S3 VPC gateway endpoint into the VPC and attach an endpoint policy that allows access to the S3 buckets.

By deploying an S3 VPC gateway endpoint, the application can access the S3 buckets over a private network connection within the VPC, eliminating the need for data transfer over the internet. This can help reduce data transfer fees as well as improve the performance of the application. The endpoint policy can be used to specify which S3 buckets the application has access to.

upvoted 47 times

Buruguduystunstugudunstuy 2 years, 5 months ago

WRONG

Option A, deploying Amazon API Gateway into a public subnet and adjusting the route table, would not address the issue of data transfer fees as the application would still be transferring data over the internet.

Option B, deploying a NAT gateway into a public subnet and attaching an endpoint policy, would not address the issue of data transfer fees either as the NAT gateway is used to enable outbound internet access for instances in a private subnet,

rather than for connecting to S3.

Option C, deploying the application into a public subnet and allowing it to route through an internet gateway, would not reduce data transfer fees as the application would still be transferring data over the internet.

upvoted 13 times

KVK16 Highly Voted 2 years, 8 months ago

Selected Answer: D

To reduce costs get rid of NAT Gateway , VPC endpoint to S3

upvoted 24 times

Gizmo2022 Most Recent 7 months, 1 week ago

D is the correct answer

upvoted 1 times

PaulGa 9 months ago

Selected Answer: D

Ans D - remove the internet connection by using a more efficient private VPC direct to S3

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

S3 VPC Gateway is the cheapest solution as it does not use any billable traffic within same region

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: D

Prevent traffic from traversing the internet = Gateway VPC endpoint for S3.

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: D

The best solution to reduce data transfer costs for an application frequently accessing S3 buckets in the same region is option D - Deploy an S3 VPC gateway endpoint into the VPC and attach an endpoint policy that allows access to the S3 buckets.

The key points:

- S3 gateway endpoints allow private connections between VPCs and S3 without going over the public internet.
- This avoids data transfer fees for traffic between the VPC and S3 within the same region.
- An endpoint policy controls access to specific S3 buckets.

upvoted 3 times

cookieMr 1 year, 11 months ago

Selected Answer: D

A. API Gateway can serve as a proxy for S3 requests, it adds unnecessary complexity and additional costs compared to a direct VPC endpoint.

B. Using a NAT gateway for accessing S3 introduces unnecessary data transfer costs as traffic would still flow over the internet.

C. This approach would incur data transfer fees as the traffic would go through the public internet.

In comparison, option D using an S3 VPC gateway endpoint provides a direct and cost-effective solution for accessing S3 buckets within the same Region. By keeping the data transfer within the AWS network infrastructure, it helps reduce data transfer fees and provides secure access to the S3 resources.

upvoted 3 times

Bmarodi 2 years ago

Selected Answer: D

Option D is correct answer.

upvoted 1 times

Erbua 2 years, 4 months ago

To answer this question, I need to know the comparison of the types of gateway of costs, please give me a tip about that issue.
upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: D

Option D

upvoted 1 times

9014 2 years, 6 months ago

Selected Answer: D

The answer is D:- Actually, the Application (EC2) is running in the same region...instead of going to the internet, data can be copied through the VPC endpoint...so there will be no cost because data is not leaving the AWS infra
upvoted 1 times

JayBee65 2 years, 6 months ago

Can somebody please explain this question? Are we assuming the application is running in AWS and that adding the gateway endpoint avoids the need for the EC2 instance to access the internet and thus avoid costs? Thanks a lot.
upvoted 2 times

SR0611 2 years, 6 months ago

Yes correct

upvoted 1 times

Wpcorgan 2 years, 6 months ago

D is correct

upvoted 1 times

yd_h 2 years, 7 months ago

Selected Answer: D

FYI :

- There is no additional charge for using gateway endpoints.
- Interface endpoints are priced at ~ \$0.01/per AZ/per hour. Cost depends on the Region
- S3 Interface Endpoints resolve to private VPC IP addresses and are routable from outside the VPC (e.g via VPN, Direct Connect, Transit Gateway, etc). S3 Gateway Endpoints use public IP ranges and are only routable from resources within the VPC.

upvoted 5 times

123jhl0 2 years, 7 months ago

Selected Answer: D

Close question to the Question #4, with same solution.

upvoted 3 times


Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions**Question #73****Topic 1**

A company recently launched Linux-based application instances on Amazon EC2 in a private subnet and launched a Linux-based bastion host on an Amazon EC2 instance in a public subnet of a VPC. A solutions architect needs to connect from the on-premises network, through the company's internet connection, to the bastion host, and to the application servers. The solutions architect must make sure that the security groups of all the EC2 instances will allow that access. Which combination of steps should the solutions architect take to meet these requirements? (Choose two.)

- A. Replace the current security group of the bastion host with one that only allows inbound access from the application instances.
- B. Replace the current security group of the bastion host with one that only allows inbound access from the internal IP range for the company.
- C. Replace the current security group of the bastion host with one that only allows inbound access from the external IP range for the company. **Most Voted**
- D. Replace the current security group of the application instances with one that allows inbound SSH access from only the private IP address of the bastion host. **Most Voted**
- E. Replace the current security group of the application instances with one that allows inbound SSH access from only the public IP address of the bastion host.

Correct Answer: CD

Community vote distribution

**Comments**

Six_Fingered_Jose **Highly Voted** 2 years, 7 months ago

Selected Answer: CD

C because from on-prem network to bastion through internet (using on-prem resource's public IP),
D because bastion and ec2 is in same VPC, meaning bastion can communicate to EC2 via it's private IP address
upvoted 50 times

slimen **Highly Voted** 1 year, 7 months ago

Selected Answer: CD

on-prem -----> bastion host (we use internet, means that we need external IPs of the company)

bastion host -----> private subnet (we use private IP since we are in the same AWS network)

upvoted 8 times

PaulGa Most Recent 9 months ago

Selected Answer: CD

Ans C,D - as per Six_Fingered_Jose (1 year, 10 mth ago)

upvoted 1 times

lofzee 1 year ago

Selected Answer: CD

CD - easy

upvoted 1 times

Anthony_Rodrigues 1 year ago

Selected Answer: CD

I've noticed that it is very important to focus on the logic for the solution, not just services.

For example, in this question, the goal is to access the application instances only from the bastion while keeping them in the private subnet, which already suggests that the SSH connection must be allowed for the bastion private IP. This is answer D.

On the other hand, the bastion must accept connections only from the company's premises, which already eliminates option A. Option B is wrong because internal IP is used only internally; in this case, the connection will be through the internet, which means that it must be the external IP; therefore, answer C.

upvoted 6 times

awsgeek75 1 year, 5 months ago

Selected Answer: CD

C: Bastion in public subnet should only allow access from public IP of the company

D: app instance in private subnet should only allow access from bastion

ABD are wrong choices

Here is a working example on AWS docs if you want to learn about Bastion setup

<https://aws.amazon.com/solutions/implementations/linux-bastion/>

upvoted 2 times

awsgeek75 1 year, 5 months ago

Good way to remember this one is to think of movie scene where someone is visiting a prisoner in a prison and talks to them from behind a glass using a 2 way phone.

Visitor is in company

Visitor area is public subnet

Prisoner area is private subnet

Phone is bastion

Visitor (company) must only be allowed to use the phone (bastion) from the public area (public subnet) and the phone (bastion) must only be allowed to talk to the prisoner in the prisoner area (private subnet)

upvoted 25 times

1dfed2b 1 year, 2 months ago

really good explanation!!!

upvoted 4 times

Marco_St 1 year, 6 months ago

Selected Answer: BD

the question mentioned from on-prem network to bastion through the company's internet then it should use the internal IP range not external ip ranges. so BD

upvoted 2 times

awsgeek75 1 year, 5 months ago

How would you know the internal IP range of the company? B is wrong.

upvoted 1 times

ATInnovandoJuntos 1 year, 6 months ago

https://en.wikipedia.org/wiki/Network_address_translation

That's the reason is C and not B

upvoted 2 times

viru 1 year, 5 months ago

Company's internet is internet - unless its company's intranet

upvoted 2 times

wearrexdzw3123 1 year, 7 months ago

Why are there always such unclear questions?

upvoted 4 times

tom_cruise 1 year, 8 months ago

Selected Answer: CD

Key: through the company's internet connection

upvoted 3 times

prabhjot 1 year, 8 months ago

Option B - inbound access from the internal IP range for the company. This step ensures that only internal IP addresses from your company's network can access the bastion host, enhancing security and then Option D

upvoted 1 times

Subhrangsu 1 year, 8 months ago

Please check first comments from top of them:

Help2023

WhericanIstart

Buruguduystunstugudunstuy

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: CD

Allows inbound access from the external IP range for the company. Then allow inbound SSH access from only the private IP address of the bastion host.

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: CD

C. This will restrict access to the bastion host from the specific IP range of the on-premises network, ensuring secure connectivity. This step ensures that only authorized users from the on-premises network can access the bastion host.

D. This step enables SSH connectivity from the bastion host to the application instances in the private subnet. By allowing inbound SSH access only from the private IP address of the bastion host, you ensure that SSH access is restricted to the bastion host only.

upvoted 3 times

stanleyjade 2 years, 1 month ago

the internal and external IP range is not clear

upvoted 4 times

PLN6302 1 year, 9 months ago

yes same for me

upvoted 1 times

pentium75 1 year, 5 months ago

There is a problem with the question. The question is asking for the correct answer, but the correct answer is not provided. The correct answer is "NOT KNOWN".

The admin is supposed to use "the company's Internet connection", NOT a VPN tunnel or DirectConnect. Thus the connection originates from the company's public/external IP.

upvoted 2 times

km142646 2 years, 1 month ago

The private/public IP address thing is confusing. Ideally, the private instances inbound rule would just allow traffic from the security group of the bastion host.

upvoted 2 times

Spiffaz 2 years, 3 months ago

Why external and not internal?

upvoted 2 times

TariqKipkemei 2 years, 3 months ago

Because the traffic goes through the public internet. In the public internet, public IP(external IP) is used.

upvoted 7 times

Help2023 2 years, 3 months ago

Selected Answer: CE

Application is in private subnet

Bastion Host is in public subnet

D does not make sense because the bastion host is in public subnet and they don't have a private IP but only a public IP address attached to them. The IP wanting to connect is Public as well.

Bastion host in public subnet allows external IP (via internet) of the company to access it. Which then leaves us to give permission to the application private subnet and for that the private subnet with the application accepts the IP coming from Bastion Host by changing its SG. C&E

upvoted 2 times

WheracanIstart 2 years, 3 months ago

Bastion host in public subnet because it has a public IP and a NAT Gateway that can route traffic out of your AWS VPC but it does have the ability to access the private subnet using private IP since it's not leaving AWS to access the private subnet. So C&D are the right answers.

upvoted 3 times

sidharthwader 1 year, 3 months ago

You are right E is also fine but it's not a best thing to do. Using private IP is always better than using public IP unless the situation demands the use of Public IP.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #74

Topic 1


A solutions architect is designing a two-tier web application. The application consists of a public-facing web tier hosted on Amazon EC2 in public subnets. The database tier consists of Microsoft SQL Server running on Amazon EC2 in a private subnet. Security is a high priority for the company.

How should security groups be configured in this situation? (Choose two.)

- A. Configure the security group for the web tier to allow inbound traffic on port 443 from 0.0.0.0/0. **Most Voted**
- B. Configure the security group for the web tier to allow outbound traffic on port 443 from 0.0.0.0/0.
- C. Configure the security group for the database tier to allow inbound traffic on port 1433 from the security group for the web tier. **Most Voted**
- D. Configure the security group for the database tier to allow outbound traffic on ports 443 and 1433 to the security group for the web tier.
- E. Configure the security group for the database tier to allow inbound traffic on ports 443 and 1433 from the security group for the web tier.

Correct Answer: AC

Community vote distribution



Comments

Athena **Highly Voted** 2 years, 7 months ago

Selected Answer: AC

Web Server Rules: Inbound traffic from 443 (HTTPS) Source 0.0.0.0/0 - Allows inbound HTTPS access from any IPv4 address
Database Rules : 1433 (MS SQL) The default port to access a Microsoft SQL Server database, for example, on an Amazon RDS instance

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/security-group-rules-reference.html>
upvoted 30 times

ArielSchivo **Highly Voted** 2 years, 7 months ago

Selected Answer: AC

EC2 web on public subnets + EC2 SQL on private subnet + security is high priority. So, Option A to allow HTTPS from

everywhere. Plus option C to allow SQL connection from the web instance.

upvoted 19 times

Dharmarajan Most Recent 4 months, 2 weeks ago

Selected Answer: AC

HTTPS (443) from the world

MS SQL server sock connection in 1433 from the private interfaces of the public web tier servers.

Security groups can be added to other security group rules. Hence A,C.

upvoted 1 times

satyaamm 5 months, 1 week ago

Selected Answer: AC

A & C are the only correct options here.

upvoted 1 times

Himedere 5 months, 3 weeks ago

Selected Answer: AC

Web server is public: should allow all (0.0.0.0/0) inbound traffic from 443 (HTTPS)

MSSQL Server is private: it should only allow inbound traffic from the web tier from 1433.

443 is HTTPS port, so it's not necessary on MSSQL Server.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: AC

Ans A, C - allow public access for all input, but control access to database: source 0.0.0.0/0, control access on port 1433 (MS SQL)

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: AC

SG are blocked by default and stateful so

A: Allows inbound traffic from web to the HTTPS default port on web servers

B: Outbound is not required if inbound is configured due to stateful nature of SG

C: 1433 is SQL default so allow access from web-tier only

D: Opens up the database to web on 1433 port

E: opens up 443 port unnecessarily on the DB tier so less secure

AC is the most secure config

upvoted 4 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: AC

Allow inbound traffic on port 443 from 0.0.0.0/0 on the web tier. Then allow inbound traffic on port 1433 from the security group for the web tier on the database tier.

upvoted 1 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: AC

The security group for the web tier should allow inbound traffic on port 443 from 0.0.0.0/0. This will allow clients to connect to the web tier using HTTPS. The security group for the web tier should also allow outbound traffic on port 443 to 0.0.0.0/0. This will allow the web tier to connect to the internet to download updates and other resources.

The security group for the database tier should allow inbound traffic on port 1433 from the security group for the web tier. This will allow the web tier to connect to the database tier to access data. The security group for the database tier should not allow outbound traffic on ports 443 and 1433 to the security group for the web tier. This will prevent the database tier from being exposed to the public internet.

upvoted 4 times

cookieMr 1 year, 11 months ago

Selected Answer: AC

A. This configuration allows external users to access the web tier over HTTPS (port 443). However, it's important to note that it is generally recommended to restrict the source IP range to a more specific range rather than allowing access from 0.0.0.0/0.

is generally recommended to restrict the source IP range to a more specific range rather than allowing access from 0.0.0.0/0 (anywhere). This would limit access to only trusted sources.

C. By allowing inbound traffic on port 1433 (default port for Microsoft SQL Server) from the security group associated with the web tier, you ensure that the database tier can only be accessed by the EC2 instances in the web tier. This provides a level of isolation and restricts direct access to the database tier from external sources.

upvoted 3 times

Abrar2022 2 years ago

DB tier: Port 1433 is the known standard for SQL server and should be used.
web tier on port 443 (HTTPS)

upvoted 4 times

beginnercloud 2 years ago

Selected Answer: AC

AC is correct

upvoted 1 times

WherecanIstart 2 years, 3 months ago

A & C are the correct answer.

Inbound traffic to the web tier on port 443 (HTTPS)
The web tier will then access the Database tier on port 1433 - inbound.

upvoted 1 times

techhb 2 years, 5 months ago

Selected Answer: AC

AC 443-http inbound and 1433-sql server
Security group => focus on inbound traffic since by default outbound traffic is allowed

upvoted 2 times

aba2s 2 years, 5 months ago

Selected Answer: AC

Security group => focus on inbound traffic since by default outbound traffic is allowed

upvoted 2 times

orionizzie 2 years, 5 months ago

why both are inbound rules

upvoted 2 times

kraken21 2 years, 2 months ago

Because security groups are stateful.

upvoted 2 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: CE

CORRECT

The correct answers are C and E.

For security purposes, it is best practice to limit inbound and outbound traffic as much as possible. In this case, the web tier should only be able to access the database tier and not the other way around. Therefore, the security group for the web tier should only allow outbound traffic to the security group for the database tier on the necessary ports. Similarly, the security group for the database tier should only allow inbound traffic from the security group for the web tier on the necessary ports.

Answer C: Configure the security group for the database tier to allow inbound traffic on port 1433 from the security group for the web tier. This is correct because the web tier needs to be able to connect to the database on port 1433 in order to access the data.

upvoted 1 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Answer E: Configure the security group for the database tier to allow inbound traffic on ports 443 and 1433 from the security

Answer E: Configure the security group for the database tier to allow inbound traffic on ports 443 and 1433 from the security group for the web tier. This is correct because the web tier needs to be able to connect to the database on both port 443 and 1433 in order to access the data.

WRONG

Answer A: Configure the security group for the web tier to allow inbound traffic on port 443 from 0.0.0.0/0. This is not correct because the web tier should not allow inbound traffic from the internet. Instead, it should only allow outbound traffic to the security group for the database tier.

upvoted 1 times

Buruguduystunstugudunstuy 2 years, 5 months ago

WRONG

Answer B: Configure the security group for the web tier to allow outbound traffic on port 443 from 0.0.0.0/0. This is not correct because the web tier should not allow outbound traffic to the internet. Instead, it should only allow outbound traffic to the security group for the database tier.

Answer D: Configure the security group for the database tier to allow outbound traffic on ports 443 and 1433 to the security group for the web tier. This is not correct because the database tier should not allow outbound traffic to the web tier. Instead, it should only allow inbound traffic from the security group for the web tier on the necessary ports.

upvoted 1 times

techhb 2 years, 5 months ago

Chatgpt is unreliable this answer from same.

upvoted 1 times

pentium75 1 year, 5 months ago

"The web tier needs to be able to connect to the database on both port 443 and 1433 in order to access the data" -> Nonsense, SQL Server needs only tcp/1433. Or tcp/1433 + udp/1433 plus the instance port if you have multiple instances. But you NEVER need tcp/443 to access SQL Server.

upvoted 1 times

PassNow1234 2 years, 5 months ago

This is WRONG. Browse to a website and type :443 at the end of it. IT will translate to HTTPS. HTTPS = 443.

answers are A and C

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #75

Topic 1

A company wants to move a multi-tiered application from on premises to the AWS Cloud to improve the application's performance. The application consists of application tiers that communicate with each other by way of RESTful services. Transactions are dropped when one tier becomes overloaded. A solutions architect must design a solution that resolves these issues and modernizes the application.

Which solution meets these requirements and is the MOST operationally efficient?

- A. Use Amazon API Gateway and direct transactions to the AWS Lambda functions as the application layer. Use Amazon Simple Queue Service (Amazon SQS) as the communication layer between application services. **Most Voted**
- B. Use Amazon CloudWatch metrics to analyze the application performance history to determine the servers' peak utilization during the performance failures. Increase the size of the application server's Amazon EC2 instances to meet the peak requirements.
- C. Use Amazon Simple Notification Service (Amazon SNS) to handle the messaging between application servers running on Amazon EC2 in an Auto Scaling group. Use Amazon CloudWatch to monitor the SNS queue length and scale up and down as required.
- D. Use Amazon Simple Queue Service (Amazon SQS) to handle the messaging between application servers running on Amazon EC2 in an Auto Scaling group. Use Amazon CloudWatch to monitor the SQS queue length and scale up when communication failures are detected.

Correct Answer: A*Community vote distribution*

A (82%)

Other (18%)

Comments**gcmrjbr** **Highly Voted** 2 years, 7 months ago

Agree with A>>> Lambda = serverless + autoscale (modernize), SQS= decouple (no more drops)
upvoted 46 times

LuckyAro **Highly Voted** 2 years, 4 months ago**Selected Answer: A**

The catch phrase is "scale up when communication failures are detected" Scaling should not be based on communication failures, that'll be crying over spilled milk ! or rather too late. So D is wrong.

upvoted 25 times

mauroicardi 1 year, 3 months ago

Spot on

upvoted 1 times

remand 2 years, 4 months ago

it says "one tier becomes overloaded" , Not communication failure...

upvoted 3 times

LuckyAro 2 years, 4 months ago

D says: "Use Amazon CloudWatch to monitor the SQS queue length and scale up when communication failures are detected".

upvoted 4 times

mc0226 Most Recent 3 months ago

Selected Answer: A

B. Analyze performance with CloudWatch and increase EC2 instance size:

Vertical scaling (increasing instance size) can address temporary performance issues but is less cost-effective and scalable. It does not modernize the architecture or address dropped transactions due to overloading.

C. Use SNS with EC2 instances and Auto Scaling:

While SNS can facilitate messaging, it is a publish/subscribe model, which is less suited for tier-to-tier decoupling. SQS is a better fit for reliable message queuing and processing.

D. Use SQS with EC2 instances and Auto Scaling:

While this addresses decoupling and scaling, it retains the burden of managing and scaling EC2 instances. A serverless solution like Lambda and API Gateway provides greater operational efficiency and modernization.

upvoted 1 times

Dharmarajan 4 months, 2 weeks ago

Selected Answer: A

Clearly A, it significantly reduces operational work. All the other options involve some maintenance and cause additional operational work to maintain the EC2 instances and ASG.

upvoted 1 times

satyaammm 5 months, 1 week ago

Selected Answer: A

Using serverless services like API Gateway and AWS Lambda and SQS helps resolve this problem with the least operational overhead.

upvoted 1 times

19d92c7 7 months, 3 weeks ago

why its not c?

upvoted 1 times

manabpokhrel7 7 months, 2 weeks ago

Operational overhead plus SNS isn't durable like SQS

upvoted 2 times

PaulGa 9 months ago

Selected Answer: A

Ans A - keep it simple: API Gateway + Lambda + SQS

D won't work: it relies upon a failure detected - by then it's too late)

upvoted 2 times

344bba0 11 months, 2 weeks ago

Selected Answer: A

The answer choice D cannot be the answer because CloudWatch is not an appropriate way to monitor transaction failures. I would have been more confused about it if CloudWatch hadn't been mentioned.

upvoted 2 times

SMALLE 1 year, 4 months ago

RESTful services = API Gateways

<https://docs.aws.amazon.com/apigateway/latest/developerguide/welcome.html>

upvoted 4 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

Least operational overhead = API Gateway + Lambda + SQS

BC wrong applications

D: Will work but more operational overhead than A with less resilience to failures

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: A

You want to "modernize the application", that a modern application is serverless, in any case a modern application does NOT use EC2 instances. Also, managing EC2 instances (with the OS etc pp) is NOT "operationally efficient". Thus not D.

upvoted 4 times

anikolov 1 year, 5 months ago

Selected Answer: C

For me the solution could be based on using SNS with multiple topics to organize communication between different tiers (Using Subscriber for one and Producer for another topic to proceed with transactions over multi-tiers). CloudWatch to monitor SNS topics queue length and scale up/down based on counts of messages (NumberOfMessagesPublished)

upvoted 1 times

MadMarc 11 months, 3 weeks ago

The question asks for the most operationally efficient solution. The answer A is serverless, while the answer C is VM-based.

There is no scenario where managin VMs is operationally more efficient than going serverless.

If Operational Efficiency is what you look for, then you should always consider going serverless.

upvoted 2 times

pentium75 1 year, 5 months ago

Wouldn't the scaling take some time? If there is queue length, then one tier IS already overloaded.

upvoted 2 times

MrPCarrot 1 year, 6 months ago

A is the perfect answer no need for the ASG

upvoted 2 times

xdkonorek2 1 year, 7 months ago

Selected Answer: D

D is better because in answer A there is a bottleneck on a SQS - service app,

D is as operationally efficient as A and solves the above issue

upvoted 2 times

tom_cruise 1 year, 8 months ago

Selected Answer: A

ASG is not as efficient as Lambda!

upvoted 2 times

vijaykamal 1 year, 8 months ago

I feel the answer is D, Lambda would increase the complexity and overhead and it has limitation of running for 15 min.

upvoted 4 times

pentium75 1 year, 5 months ago

How would Lambda increase the operational (!) complexity over handling and scaling servers (and their operating systems and natches etc.) on EC2?

practice story on ECR.

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: A

MOST operationally efficient = Serverless = AWS Lambda functions, Amazon Simple Queue Service

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #76

Topic 1

A company receives 10 TB of instrumentation data each day from several machines located at a single factory. The data consists of JSON files stored on a storage area network (SAN) in an on-premises data center located within the factory. The company wants to send this data to Amazon S3 where it can be accessed by several additional systems that provide critical near-real-time analytics. A secure transfer is important because the data is considered sensitive.

Which solution offers the MOST reliable data transfer?

- A. AWS DataSync over public internet
- B. AWS DataSync over AWS Direct Connect **Most Voted**
- C. AWS Database Migration Service (AWS DMS) over public internet
- D. AWS Database Migration Service (AWS DMS) over AWS Direct Connect

Correct Answer: B

Community vote distribution

B (100%)

Comments

ArielSchivo **Highly Voted** 2 years, 7 months ago

Selected Answer: B

DMS is for databases and here refers to "JSON files". Public internet is not reliable. So best option is B.
upvoted 38 times

Buruguduystunstugudunstuy **Highly Voted** 2 years, 5 months ago

Selected Answer: B

CORRECT

The most reliable solution for transferring the data in a secure manner would be option B: AWS DataSync over AWS Direct Connect.

AWS DataSync is a data transfer service that uses network optimization techniques to transfer data efficiently and securely between on-premises storage systems and Amazon S3 or other storage targets. When used over AWS Direct Connect, DataSync can provide a dedicated and secure network connection between your on-premises data center and AWS. This can help to ensure a more reliable and secure data transfer compared to using the public internet.

upvoted 21 times

Buruguduystunstugudunstuy 2 years, 5 months ago

ParagdeepJain 1 year, 6 months ago

WRONG

Option A, AWS DataSync over the public internet, is not as reliable as using Direct Connect, as it can be subject to potential network issues or congestion.

Option C, AWS Database Migration Service (DMS) over the public internet, is not a suitable solution for transferring large amounts of data, as it is designed for migrating databases rather than transferring large amounts of data from a storage area network (SAN).

Option D, AWS DMS over AWS Direct Connect, is also not a suitable solution, as it is designed for migrating databases and may not be efficient for transferring large amounts of data from a SAN.

upvoted 14 times

doorahmie 2 years, 4 months ago

explanation about D option is good

upvoted 1 times

satyaamm Most Recent 5 months, 1 week ago

Selected Answer: B

AWS DataSync supports Encryption and Compression and it is much more safer over a Direct Connect.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: B

Ans B - keep it simple... private netwk + direct connection

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

near-real-time + large data + secure = DataSync over DirectConnect

A: Less secure due to public internet

C: Slow and not secure

D: Slow even if more secure

DC may not even work as we don't know if there is a DB on other side but even if it was there, it is less preferred way

upvoted 3 times

JTruong 1 year, 5 months ago

Any DMS related-service will not be efficient because DMS can only process JSON files UPTO 2 MB in size
[https://docs.aws.amazon.com/dms/latest/userguide/CHAP_Tasks.CustomizingTasks.TableMapping.SelectionTransformation.htm](https://docs.aws.amazon.com/dms/latest/userguide/CHAP_Tasks.CustomizingTasks.TableMapping.SelectionTransformation.html)
l

so B is CORRECT

upvoted 2 times

Ruffyit 1 year, 7 months ago

AWS DataSync is a data transfer service that uses network optimization techniques to transfer data efficiently and securely between on-premises storage systems and Amazon S3 or other storage targets. When used over AWS Direct Connect, DataSync can provide a dedicated and secure network connection between your on-premises data center and AWS. This can help to ensure a more reliable and secure data transfer compared to using the public internet.

upvoted 4 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: B

Secure and Most reliable transfer = AWS DataSync over AWS Direct Connect

upvoted 3 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: B

AWS DataSync is designed for large scale, high speed data transfer between on-prem and S3.

Using AWS Direct Connect provides a dedicated, private connection for reliable, consistent data transfer.

DataSync seamlessly handles data replication, encryption, recovery etc.

upvoted 3 times

MNotABot 1 year, 11 months ago

Not over public hence AC out / DMS is for databases and here refers to "JSON files".

upvoted 1 times

cookieMr 1 year, 11 months ago

Selected Answer: B

DataSync is a service specifically designed for data transfer and synchronization between on-premises storage systems and AWS storage services like S3. It provides reliable and efficient data transfer capabilities, ensuring the secure movement of large volumes of data.

By leveraging Direct Connect, which establishes a dedicated network connection between the on-premises data center and AWS, the data transfer is conducted over a private and dedicated network link. This approach offers increased reliability, lower latency, and consistent network performance compared to transferring data over the public internet.

Database Migration Service is primarily focused on database migration and replication, and it may not be the most appropriate tool for general-purpose data transfer like JSON files.

Transferring data over the public internet may introduce potential security risks and performance variability due to factors like network congestion, latency, and potential interruptions.

upvoted 3 times

beginnercloud 2 years ago

Best option and correct is B

upvoted 1 times

Abrar2022 2 years ago

Selected Answer: B

as Ariel suggested and rightly so.....DMS is for databases and here refers to "JSON files". Public internet is not reliable. so B

upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: B

Option B

upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: B

Option B. DMS is not needed as there is no Database migration requirement.

upvoted 1 times

Wajif 2 years, 6 months ago

Selected Answer: B

Public internet options automatically out being best-effort. DMS is for database migration service and here they have to just transfer the data to S3. Hence B.

upvoted 2 times

Wpcorgan 2 years, 6 months ago

B is correct

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #77

Topic 1


A company needs to configure a real-time data ingestion architecture for its application. The company needs an API, a process that transforms data as the data is streamed, and a storage solution for the data.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Deploy an Amazon EC2 instance to host an API that sends data to an Amazon Kinesis data stream. Create an Amazon Kinesis Data Firehose delivery stream that uses the Kinesis data stream as a data source. Use AWS Lambda functions to transform the data. Use the Kinesis Data Firehose delivery stream to send the data to Amazon S3.
- B. Deploy an Amazon EC2 instance to host an API that sends data to AWS Glue. Stop source/destination checking on the EC2 instance. Use AWS Glue to transform the data and to send the data to Amazon S3.
- C. Configure an Amazon API Gateway API to send data to an Amazon Kinesis data stream. Create an Amazon Kinesis Data Firehose delivery stream that uses the Kinesis data stream as a data source. Use AWS Lambda functions to transform the data. Use the Kinesis Data Firehose delivery stream to send the data to Amazon S3. **Most Voted**
- D. Configure an Amazon API Gateway API to send data to AWS Glue. Use AWS Lambda functions to transform the data. Use AWS Glue to send the data to Amazon S3.

Correct Answer: C

Community vote distribution



Comments

123jh10 **Highly Voted** 2 years, 7 months ago

Selected Answer: C

(A) - You don't need to deploy an EC2 instance to host an API - Operational overhead

(B) - Same as A

(**C**) - Is the answer

(D) - AWS Glue gets data from S3, not from API GW. AWS Glue could do ETL by itself, so don't need lambda. Non sense.

<https://aws.amazon.com/glue/>

upvoted 44 times

Futurebones 2 years, 1 month ago

I don't understand is why we should use Lambda in between to transform data. To me, Kinesis data firehose is enough as it is an extract, transform, and load (ETL) service.

upvoted 5 times

Remy_d 1 year, 8 months ago

It is because they assume that Kinesis Data Firehose built-in transformations are not enough. So you have to use specific lambda transformation. Please refer to this link : <https://aws.amazon.com/kinesis/data-firehose/#:~:text=Amazon%20Kinesis%20Data%20Firehose%20is,data%20stores%2C%20and%20analytics%20services.>

upvoted 9 times

TariqKipkemei Highly Voted 1 year, 10 months ago

Selected Answer: C

The company needs an API = Amazon API Gateway API
A real-time data ingestion = Amazon Kinesis data stream
A process that transforms data = AWS Lambda functions
Kinesis Data Firehose delivery stream to send the data to Amazon S3
A storage solution for the data = Amazon S3

upvoted 30 times

PaulGa Most Recent 9 months ago

Selected Answer: C

Ans C - the API is given so just configure Kinesis stream/Firehose to use it

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

C is least operational overhead

A: EC2 is overhead in this scenario
B: Same as A
D: Glue is not real time data streaming

upvoted 3 times

Mikado211 1 year, 6 months ago

Selected Answer: C

It looks overengineered, but as it works, let's go for the C

upvoted 5 times

Ruffyt 1 year, 7 months ago

The company needs an API = Amazon API Gateway API
A real-time data ingestion = Amazon Kinesis data stream
A process that transforms data = AWS Lambda functions
Kinesis Data Firehose delivery stream to send the data to Amazon S3
A storage solution for the data = Amazon S3

upvoted 4 times

peekingpicker 1 year, 8 months ago

Selected Answer: D

"a real-time data ingestion"

isn't firehose not realtime ? Kinesis FireHose is "Near" Real-time . It has 60 seconds gap. I think it should be D

upvoted 1 times

rlamberti 1 year, 7 months ago

The real-time part (data ingestion) will be performed by Kinesis Data Stream and API Gateway. After this, the transformation and storage of the data don't need to be in real-time, since it was already ingested, so Kinesis Firehose + Lambda is perfect. C makes sense to me.

upvoted 5 times

Guru4Cloud 1 year, 10 months ago

Selected Answer: C

Option C provides the least operational overhead to meet the requirements:

API Gateway provides the API
Kinesis Data Streams ingests the real-time data
Lambda functions transform the data
Firehose delivers the data to S3 storage
The key advantages are:

Serverless architecture requires minimal operational overhead
Fully managed ingestion, processing and storage services
No need to manage EC2 instances
upvoted 2 times

diabloexodia 1 year, 11 months ago

Requirements:
API- API gateway
Real time data ingestion - AWS Kinesis data stream
ETL(Extract Transform Load) - Kinesis Firehose
Storage- S3
upvoted 3 times

tamefi5512 1 year, 11 months ago

Selected Answer: C

C - is the answer
upvoted 1 times

cookieMr 1 year, 11 months ago

Selected Answer: C

C. By leveraging these services together, you can achieve a real-time data ingestion architecture with minimal operational overhead. The data flows from the API Gateway to the Kinesis data stream, undergoes transformations with Lambda, and is then sent to S3 via the Kinesis Data Firehose delivery stream for storage.

A. This adds operational overhead as you need to handle EC2 management, scaling, and maintenance. It is less efficient compared to using a serverless solution like API Gateway.

B. It requires deploying and managing an EC2 to host the API and configuring Glue. This adds operational overhead, including EC2 management and potential scalability limitations.

D. It still requires managing and configuring Glue, which adds operational overhead. Additionally, it may not be the most efficient solution as Glue is primarily used for ETL scenarios, and in this case, real-time data transformation is required.

upvoted 2 times

winzzhhzzhh 2 years ago

Selected Answer: D

I am gonna choose D for this.
Kinesis Data Stream + Data Firehose will adds up to the operational overhead, plus it is "Near real-time", not a real time solution.
Lambda functions scale automatically, so no management of scaling/compute resources is needed.
AWS Glue handles the data storage in S3, so no management of that is needed.
upvoted 2 times

UnluckyDucky 2 years, 2 months ago

Gotta love all those chatgpt answers y'all are throwing at us.

Kinesis Firehose is NEAR real-time, not real-time like your bots tell you.
upvoted 2 times

pentium75 1 year, 5 months ago

Stem is about "real-time data INGESTION", not real-time processing.
upvoted 2 times

bullrem 2 years, 4 months ago

Selected Answer: C

option C is the best solution. It uses Amazon Kinesis Data Firehose which is a fully managed service for delivering real-time

streaming data to destinations such as Amazon S3. This service requires less operational overhead as compared to option A, B, and D. Additionally, it also uses Amazon API Gateway which is a fully managed service for creating, deploying, and managing APIs. These services help in reducing the operational overhead and automating the data ingestion process.

upvoted 1 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Selected Answer: C

Option C is the solution that meets the requirements with the least operational overhead.

In Option C, you can use Amazon API Gateway as a fully managed service to create, publish, maintain, monitor, and secure APIs. This means that you don't have to worry about the operational overhead of deploying and maintaining an EC2 instance to host the API.

Option C also uses Amazon Kinesis Data Firehose, which is a fully managed service for delivering real-time streaming data to destinations such as Amazon S3. With Kinesis Data Firehose, you don't have to worry about the operational overhead of setting up and maintaining a data ingestion infrastructure.

upvoted 1 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Finally, Option C uses AWS Lambda, which is a fully managed service for running code in response to events. With AWS Lambda, you don't have to worry about the operational overhead of setting up and maintaining a server to run the data transformation code.

Overall, Option C provides a fully managed solution for real-time data ingestion with minimal operational overhead.

upvoted 2 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option A is incorrect because it involves deploying an EC2 instance to host an API, which adds operational overhead in the form of maintaining and securing the instance.

Option B is incorrect because it involves deploying an EC2 instance to host an API and disabling source/destination checking on the instance. Disabling source/destination checking can make the instance vulnerable to attacks, which adds operational overhead in the form of securing the instance.

upvoted 2 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option D is incorrect because it involves using AWS Glue to send the data to Amazon S3, which adds operational overhead in the form of maintaining and securing the AWS Glue infrastructure.

Overall, Option C is the best choice because it uses fully managed services for the API, data transformation, and data delivery, which minimizes operational overhead.

upvoted 2 times

career360guru 2 years, 5 months ago

Selected Answer: C

Option C

upvoted 1 times

career360guru 2 years, 5 months ago

Selected Answer: C

Option C

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions**Question #78**

Topic 1

A company needs to keep user transaction data in an Amazon DynamoDB table. The company must retain the data for 7 years. What is the MOST operationally efficient solution that meets these requirements?

- A. Use DynamoDB point-in-time recovery to back up the table continuously.
- B. Use AWS Backup to create backup schedules and retention policies for the table. **Most Voted**
- C. Create an on-demand backup of the table by using the DynamoDB console. Store the backup in an Amazon S3 bucket. Set an S3 Lifecycle configuration for the S3 bucket.
- D. Create an Amazon EventBridge (Amazon CloudWatch Events) rule to invoke an AWS Lambda function. Configure the Lambda function to back up the table and to store the backup in an Amazon S3 bucket. Set an S3 Lifecycle configuration for the S3 bucket.

Correct Answer: B*Community vote distribution***B (100%)****Comments****123jhl0** **Highly Voted** 2 years, 7 months ago**Selected Answer: B**

Answer is B

"Amazon DynamoDB offers two types of backups: point-in-time recovery (PITR) and on-demand backups. (==> D is not the answer)

PITR is used to recover your table to any point in time in a rolling 35 day window, which is used to help customers mitigate accidental deletes or writes to their tables from bad code, malicious access, or user error. (==> A isn't the answer)

On demand backups are designed for long-term archiving and retention, which is typically used to help customers meet compliance and regulatory requirements.

This is the second of a series of two blog posts about using AWS Backup to set up scheduled on-demand backups for Amazon DynamoDB. Part 1 presents the steps to set up a scheduled backup for DynamoDB tables from the AWS Management Console." (==> Not the DynamoDB console and C isn't the answer either)

<https://aws.amazon.com/blogs/database/part-2-set-up-scheduled-backups-for-amazon-dynamodb-using-aws-backup/>
upvoted 52 times

LuckyAro 2 years, 4 months ago

I think the answer is C because of storage time.

upvoted 2 times

jerryl 2 months, 3 weeks ago

i think another reason for why C is not correct is the keyword "on-demand" which is not operation efficient (as this should be an ongoing practice)

upvoted 1 times

MutiverseAgent 1 year, 11 months ago

Dynamo backups cannot be scheduled or sent to S3, so answer should be B)

1) <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/BackupRestore.html>

2) <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Backup.Tutorial.html>

upvoted 2 times

app12 1 year, 5 months ago

In the very same link you shared it says that you CAN send backups to S3

https://youtu.be/4lNEu_hw30Q?t=54

upvoted 2 times

Buruguduystunstugudunstuy Highly Voted 2 years, 5 months ago

Selected Answer: B

The most operationally efficient solution that meets these requirements would be to use option B, which is to use AWS Backup to create backup schedules and retention policies for the table.

AWS Backup is a fully managed backup service that makes it easy to centralize and automate the backup of data across AWS resources. It allows you to create backup policies and schedules to automatically back up your DynamoDB tables on a regular basis. You can also specify retention policies to ensure that your backups are retained for the required period of time. This solution is fully automated and requires minimal maintenance, making it the most operationally efficient option.

upvoted 18 times

Buruguduystunstugudunstuy 2 years, 5 months ago

Option A, using DynamoDB point-in-time recovery, is also a viable option but it requires continuous backup, which may be more resource-intensive and may incur higher costs compared to using AWS Backup.

Option C, creating an on-demand backup of the table and storing it in an S3 bucket, is also a viable option but it requires manual intervention and does not provide the automation and scheduling capabilities of AWS Backup.

Option D, using Amazon EventBridge (CloudWatch Events) and a Lambda function to back up the table and store it in an S3 bucket, is also a viable option but it requires more complex setup and maintenance compared to using AWS Backup.

upvoted 15 times

Dharmarajan Most Recent 4 months, 2 weeks ago

Selected Answer: B

Anytime the "Operational efficiency" comes into picture, one should consider the options with serverless or managed services. DynamoDB is a serverless service, and AWS Backup is a managed service. Others involve operational work in terms of maintenance of backups etc.

upvoted 2 times

satyaammm 5 months, 1 week ago

Selected Answer: B

Using AWS Backup is most suitable here for retaining data for 7 years.

upvoted 1 times

PaulGa 9 months ago

Selected Answer: B

Ans B - there are no special parameters specified for the backup (eg. PITR snapshots) so it doesn't need to be elaborate

upvoted 2 times

OctavioBatera 1 year, 2 months ago

Selected Answer: B

Agreed with option B is the right one. AWS backup retention goes from 1 day to 100 years (or even indefinitely, if you do not