

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #239

Topic 1

A solutions architect needs to design a new microservice for a company's application. Clients must be able to call an HTTPS endpoint to reach the microservice. The microservice also must use AWS Identity and Access Management (IAM) to authenticate calls. The solutions architect will write the logic for this microservice by using a single AWS Lambda function that is written in Go 1.x.

Which solution will deploy the function in the MOST operationally efficient way?

- A. Create an Amazon API Gateway REST API. Configure the method to use the Lambda function. Enable IAM authentication on the API. **Most Voted**
- B. Create a Lambda function URL for the function. Specify AWS_IAM as the authentication type.
- C. Create an Amazon CloudFront distribution. Deploy the function to Lambda@Edge. Integrate IAM authentication logic into the Lambda@Edge function.
- D. Create an Amazon CloudFront distribution. Deploy the function to CloudFront Functions. Specify AWS_IAM as the authentication type.

Correct Answer: A*Community vote distribution*

A (54%)

B (46%)

Comments**mhmt4438** **Highly Voted** 2 years, 4 months ago**Selected Answer: A**

A. Create an Amazon API Gateway REST API. Configure the method to use the Lambda function. Enable IAM authentication on the API.

This option is the most operationally efficient as it allows you to use API Gateway to handle the HTTPS endpoint and also allows you to use IAM to authenticate the calls to the microservice. API Gateway also provides many additional features such as caching, throttling, and monitoring, which can be useful for a microservice.

upvoted 22 times

FlyingHawk 4 months, 2 weeks ago

Simplicity:

Lambda Function URLs provide a direct HTTPS endpoint with minimal configuration. You don't need to set up API Gateway, which reduces complexity.

IAM Authentication:

Lambda Function URLs natively support IAM authentication, eliminating the need for custom logic or additional services.
upvoted 2 times

OlehKom Highly Voted 9 months, 1 week ago

The question specifically asks for the solution that is most operationally efficient, not necessarily the one with the most features (which is A).

Option B—using a Lambda function URL with AWS_IAM authentication—is indeed the most operationally efficient because:

- 1) Minimal Configuration: Lambda function URLs are designed to quickly create HTTPS endpoints without the need for additional AWS services like API Gateway.
- 2) Built-in IAM Authentication: You can easily specify AWS_IAM as the authentication type directly, fulfilling the requirement without extra setup.
- 3) No Extra Overhead: Unlike API Gateway, there's no need to manage complex API configurations, throttling settings, or additional API management features unless specifically required.

upvoted 6 times

Yak_Yeti Most Recent 1 month, 3 weeks ago

Selected Answer: B

Lambda function URLs with AWS_IAM authentication provide the most streamlined and operationally efficient way to deploy a single Go Lambda function behind an HTTPS endpoint with IAM-based authentication for client calls. It minimizes the number of moving parts and configuration required, leading to easier management and lower potential for operational issues.

upvoted 1 times

FlyingHawk 6 months, 2 weeks ago

Selected Answer: B

Use Lambda Function URLs for simplicity, low cost, and direct integration with IAM for authentication. If your microservice evolves to require more complex features, you can later integrate with API Gateway as needed. API Gateway can integrate with Lambda and enable IAM authentication, but it introduces more operational overhead and cost compared to a Lambda Function URL. It's unnecessary unless the application requires additional features like request validation, throttling, or custom authorizers.

upvoted 3 times

LeonSauveterre 6 months, 2 weeks ago

Selected Answer: B

Why is B better than A: Both options meet the requirements, but Lambda function URLs are simpler and involve less operational overhead compared to setting up and managing an API Gateway, so B for sure.

upvoted 3 times

maryam_sh 10 months ago

Selected Answer: B

B is the most operationally efficient solution. It provides the necessary functionality with minimal setup and cost, directly supports IAM authentication, and avoids the additional complexity and overhead associated with other options. This approach is particularly suitable for scenarios where the API requirements are straightforward and do not need the advanced features provided by API Gateway or CloudFront.

upvoted 4 times

hro 1 year, 2 months ago

I think from a decoupling and separation of concerns A is the answer. You don't want to have a heavy reliance on the Lambda function with you have specific services for what is being required.

there is operationally efficient incorrect and operationally efficient correct.

So A is the best answer.

upvoted 2 times

bujuman 1 year, 3 months ago

Selected Answer: B

According to this statement "MOST operationally efficient way" and the following link related to Lambda Function URL security:

<https://docs.aws.amazon.com/lambda/latest/dg/urls-auth.html>

upvoted 3 times

awsgeek75 1 year, 4 months ago

I originally voted B but after reading this article, I am not sure if A is wrong or is just badly worded. If A actually said "Configure the [authorization] method to use the Lambda function" then it would be way more logical than B but this could be intentional. Although I think this is AWS test not IELTS so picking right answers based on small word mistakes is not the intention!

<https://docs.aws.amazon.com/apigateway/latest/developerguide/http-api-lambda-authorizer.html>

upvoted 3 times

upliftinghut 1 year, 5 months ago

Selected Answer: B

A& B are all good. The requirement is most operationally efficient so B is faster. In real life, I won't risk B for production, dev & test makes sense but no production please

upvoted 2 times

upliftinghut 1 year, 5 months ago

Reference link why B is here: <https://aws.amazon.com/blogs/aws/announcing-aws-lambda-function-urls-built-in-https-endpoints-for-single-function-microservices/>

upvoted 4 times

awsgeek75 1 year, 5 months ago

Selected Answer: B

I think this question has 2 answers as both A and B will work. However, B is more operationally efficient due to Lambda function URL and direct support for AWS_IAM as the auth type for this setup.

<https://docs.aws.amazon.com/lambda/latest/dg/urls-auth.html#urls-auth-iam>

C, D are not operationally efficient and GO is not supported on Lambda@Edge or CloudFront functions. Even if AWS start supporting it, the operational efficiency will increase because of CloudFront

upvoted 2 times

awsgeek75 1 year, 5 months ago

* meant to say "operational efficiency will decrease" because of CloudFront

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: B

We know that the application provides "an HTTPS endpoint" but we don't even know whether it is a REST API. The question is not mentioning any other requirements besides IAM authentication, which can be handled by Lambda alone.

A would work, but would be an additional processing step (lowering operational efficiency). It would also provide benefits but none of those is asked for in the question.

C and D is wrong because Lambda@Edge does not support Go.

upvoted 3 times

meowruki 1 year, 6 months ago

Selected Answer: A

Options B, C, and D involve using Lambda function URLs or CloudFront, but they lack the full set of features provided by API Gateway, such as built-in IAM authentication, throttling, and other API management capabilities.

upvoted 3 times

google_platform_team 1 year, 6 months ago

Selected Answer: B

I think it is B - most operationally efficient. A is a better answer, but more complicated.

upvoted 2 times

swap001 1 year, 7 months ago

Selected Answer: B

There is no need of an additional API gateway when Lambda itself can support the need. This is more operationally efficient.
upvoted 3 times

OlehKom 1 year, 8 months ago

Why not B? I agree that A is a nice choice, but it clearly says "MOST operationally efficient way", there is nothing said about API. B in this case suits absolutely fine, it's simpler and cheaper.
upvoted 2 times

LeonSauveterre 6 months, 2 weeks ago

I totally agree. The answer should be B instead of A. I'm sure of it.
upvoted 1 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: A

Create an Amazon API Gateway REST API. Configure the method to use the Lambda function. Enable IAM authentication on the API
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #240

Topic 1


A company previously migrated its data warehouse solution to AWS. The company also has an AWS Direct Connect connection. Corporate office users query the data warehouse using a visualization tool. The average size of a query returned by the data warehouse is 50 MB and each webpage sent by the visualization tool is approximately 500 KB. Result sets returned by the data warehouse are not cached.

Which solution provides the LOWEST data transfer egress cost for the company?

- A. Host the visualization tool on premises and query the data warehouse directly over the internet.
- B. Host the visualization tool in the same AWS Region as the data warehouse. Access it over the internet.
- C. Host the visualization tool on premises and query the data warehouse directly over a Direct Connect connection at a location in the same AWS Region.
- D. Host the visualization tool in the same AWS Region as the data warehouse and access it over a Direct Connect connection at a location in the same Region. **Most Voted**

Correct Answer: D

Community vote distribution

**Comments**

AlessandraSAA **Highly Voted** 2 years, 3 months ago

Selected Answer: D

- A. --> No since if you access via internet you are creating egress traffic.
- B. --> It's a good choice to have both DWH and visualization in the same region to lower the egress transfer (i.e. data going egress/out of the region) but if you access over internet you might still have egress transfer.
- C. -> Valid but in this case you send out of AWS 50MB if you query the DWH instead of the visualization tool, D removes this need since puts the visualization tools in AWS with the DWH so reduces data returned out of AWS from 50MB to 500KB
- D. --> Correct, see explanation on answer C

Useful links:

AWS Direct Connect connection create a connection in an AWS Direct Connect location to establish a network connection from your premises to an AWS Region.

<https://docs.aws.amazon.com/directconnect/latest/UserGuide/Welcome.html>

upvoted 15 times

Jacky_S 11 months, 3 weeks ago

It is not quite making sense that the answer D to mention on AWS Direct Connect if the tool is staying with AWS already. That is the main reason i vote on C rather than D.

upvoted 3 times

hro 1 year, 2 months ago

I thought Direct Connect is exclusively for on-prem to VPC. Why would you use DirectConnect to connect to a service in your own Region in the AWS Cloud?

I believe C is correct.

upvoted 3 times

chickenmf 1 year, 2 months ago

Agree with you 100%

upvoted 1 times

cookieMr Highly Voted 1 year, 11 months ago

Selected Answer: D

Hosting the visualization tool in the same AWS Region as the data warehouse and accessing it over a Direct Connect connection within the same Region eliminates data transfer fees and ensures low-latency, high-bandwidth connectivity.

A. Hosting the visualization tool on premises and querying the data warehouse over the internet incurs data transfer costs for every query result, as well as potential latency and bandwidth limitations.

B. Hosting the visualization tool in the same AWS Region as the data warehouse but accessing it over the internet still incurs data transfer costs for each query result.

C. Hosting the visualization tool on premises and querying the data warehouse over a Direct Connect connection within the same AWS Region incurs data transfer costs for every query result and adds complexity by requiring on-premises infrastructure.

upvoted 5 times

LeonSauveterre Most Recent 6 months, 2 weeks ago

Selected Answer: D

Why C is Less Cost-Effective Than D: While Direct Connect minimizes costs in both cases, hosting the visualization tool on premises (C) still incurs some egress costs when pulling large query results to the corporate office. Hosting the visualization tool in the same AWS Region as the data warehouse (D) eliminates most egress costs ****entirely**** due to AWS's intra-region free data transfer policy.

upvoted 2 times

Jazz888 1 year ago

This is question that need a careful reading of the question. Please note: in the question that the company has a DX connection but it doesn't mention the company is utilizing the DX to query the data warehouse (this indicates that the query was using internet). So my verdict would be to place the visualization tool on-prem (which it already is) and use the DX to query the data warehouse to reduce cost of (DTO) so I vote for C

upvoted 2 times

hro 1 year, 2 months ago

The answer is C. At no point does the question suggest that the DWH source is out of Region.

upvoted 2 times

reviewmine 1 year, 3 months ago

Selected Answer: D

It's D. Host the visualization on the same region to avoid egress cost and access the tool via AWS Direct connection.

upvoted 3 times

upliftinghut 1 year, 5 months ago

Selected Answer: D

Leverage the existing DirectConnect so not incur data transfer charge

upvoted 2 times

Ruffyt 1 year, 6 months ago

D. Host the visualization tool in the same AWS Region as the data warehouse and access it over a Direct Connect connection at a location in the same Region.

upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: D

Host the visualization tool in the same AWS Region as the data warehouse and access it over a Direct Connect connection at a location in the same Region

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

D. Host the visualization tool in the same AWS Region as the data warehouse and access it over a Direct Connect connection at a location in the same Region.

upvoted 2 times

jtexam 1 year, 11 months ago

Selected Answer: B

by hosting in same region, you have 500kb transfer charged on internet transfer tier, 50MB charged in inter-region tier.

using direct link, both are charged in direct link tier. direct link tier is not cheap.

so i go for B

upvoted 1 times

pentium75 1 year, 5 months ago

We don't want "lowest cost", we want "lowest data transfer egress cost". And "data transfer egress" cost for Direct Connect is WAY lower than for Internet. (Therefore the circuit itself is expensive, but that does exist anyway.)

upvoted 2 times

Mmmmmmkkkk 1 year, 11 months ago

Aaaaaaaa

upvoted 1 times

dexpos 2 years, 4 months ago

Selected Answer: D

D let you reduce at minimum the data transfer costs

upvoted 2 times

alexleely 2 years, 4 months ago

Selected Answer: D

D: Direct Connect connection at a location in the same Region will provide the lowest data transfer egress cost, improved performance, and lower complexity

Why it is not C is because the visualization tool is hosted on-premises, as it's not hosted in the same region as the data warehouse the data transfer between them would occur over the internet, thus, would incur in egress data transfer costs.

upvoted 5 times

markw92 1 year, 11 months ago

C option doesn't travel through internet because we have a direct connect. If you are hosting your visualization tool in same region why you need a direct connection which D has? Doesn't make sense. So, C is the right answer.

upvoted 1 times

Vickysss 2 years, 4 months ago

Selected Answer: C

<https://www.nops.io/reduce-aws-data-transfer-costs-dont-get-stung-by-hefty-egress-fees/>

upvoted 2 times

JayBee65 2 years, 4 months ago

Whilst "Direct Connect can help lower egress costs even after taking the installation costs into account. This is because AWS charges lower transfer rates." D removes the need to send the query results out of AWS and instead returns the web page, so reduces data returned from 50MB to 500KB, so D

upvoted 3 times

mhmt4438 2 years, 4 months ago

Selected Answer: D

Correct answer is D

upvoted 5 times

Aninina 2 years, 4 months ago

Selected Answer: D

Should be D

<https://aws.amazon.com/directconnect/pricing/>

<https://aws.amazon.com/blogs/aws/aws-data-transfer-prices-reduced/>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #241

Topic 1

An online learning company is migrating to the AWS Cloud. The company maintains its student records in a PostgreSQL database. The company needs a solution in which its data is available and online across multiple AWS Regions at all times.

Which solution will meet these requirements with the LEAST amount of operational overhead?

- A. Migrate the PostgreSQL database to a PostgreSQL cluster on Amazon EC2 instances.
- B. Migrate the PostgreSQL database to an Amazon RDS for PostgreSQL DB instance with the Multi-AZ feature turned on.
- C. Migrate the PostgreSQL database to an Amazon RDS for PostgreSQL DB instance. Create a read replica in another Region. **Most Voted**
- D. Migrate the PostgreSQL database to an Amazon RDS for PostgreSQL DB instance. Set up DB snapshots to be copied to another Region.

Correct Answer: C*Community vote distribution***Comments****Steve_4542636** **Highly Voted** 2 years, 3 months ago**Selected Answer: C**

Multi az is not the same as multi regional
upvoted 48 times

alexleely **Highly Voted** 2 years, 4 months ago**Selected Answer: B**

B: Amazon RDS Multi-AZ feature automatically creates a synchronous replica in another availability zone and failover to the replica in the event of an outage. This will provide high availability and data durability across multiple AWS regions which fit the requirements.

Though C may sound good, it in fact requires manual management and monitoring of the replication process due to the fact that Amazon RDS read replicas are asynchronous, meaning there is a delay between the primary and read replica. Therefore, there will be a need to ensure that the read replica is constantly up-to-date and someone still has to fix any read replica errors during the replication process which may cause data inconsistency. Lastly, you still have to configure additional steps to make it fail over to the read replica.

upvoted 16 times

Rehan33 2 years, 3 months ago

I go with option B because:
Multi-AZ is for high availability
Read replicas are for low-latency
in question they talk about available online
upvoted 4 times

awsgeek75 1 year, 4 months ago

They also ask for multiple regions which is not covered by Multi AZ
upvoted 3 times

Mahadeva 2 years, 4 months ago

But the question is clearly asking for Multiple Regions. Multi-AZ is not across Regions.
upvoted 22 times

alexleely 2 years, 4 months ago

You are right, Multi-AZ is only within one Region. C would be the right answer.
upvoted 15 times

smartegnine 1 year, 12 months ago

<https://aws.amazon.com/rds/features/multi-az/>

smartegnine 0 minutes ago Awaiting moderator approval
Selected Answer: B
In an Amazon RDS Multi-AZ deployment, Amazon RDS automatically creates a primary database (DB) instance and synchronously replicates the data to an instance in a different AZ.
upvoted 1 times

hasiet Most Recent 6 months, 1 week ago

Selected Answer: C

C. Migrate the PostgreSQL database to an Amazon RDS for PostgreSQL DB instance. Create a read replica in another Region.

This option provides cross-region availability with minimal operational overhead, as RDS manages the replication process for you.
upvoted 2 times

akshay243007 10 months, 2 weeks ago

Selected Answer: C

Multi AZ is not supported in to different region
read Replica is supported in to different region

So I am go with the option C
upvoted 2 times

Uzbekistan 1 year, 2 months ago

Selected Answer: B

Option C involves setting up read replicas in another region, which provides cross-region availability but may introduce additional complexity in managing replication and monitoring.

B. Migrate the PostgreSQL database to an Amazon RDS for PostgreSQL DB instance with Multi-AZ:

High Availability: Amazon RDS Multi-AZ deployments provide high availability and failover support for DB instances. With Multi-AZ, Amazon RDS automatically provisions and maintains a synchronous standby replica in a different Availability Zone (AZ). In the event of a failure, Amazon RDS automatically fails over to the standby replica, ensuring data availability and minimal downtime. + Cross-Region Replication.
upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: C

- A - Just no
 - B - Multi-AZ is multiple AZs in same region, does not meet "Multiple AWS Regions" requirement
 - C - Meets requirements
 - D - Does not meet the "online ... at all times" requirement
- upvoted 10 times

Ruffyt 1 year, 6 months ago

Multi az is not the same as multi regional

upvoted 2 times

vijaykamal 1 year, 8 months ago

Selected Answer: B

Option C, while providing a read replica in another Region, adds complexity to the architecture and may introduce some additional operational overhead compared to Multi-AZ. Cross-Region replication involves setting up and managing replication between two separate RDS instances.

upvoted 1 times

pentium75 1 year, 5 months ago

But data must be online "across multiple AWS regions"!

upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: C

Migrate the PostgreSQL database to an Amazon RDS for PostgreSQL DB instance. Create a read replica in another Region

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

Multi-AZ is not the same as Multi-Regional

upvoted 4 times

Valder21 1 year, 9 months ago

can someone explain why not D

upvoted 3 times

pentium75 1 year, 5 months ago

"Data must be online (!) across multiple AWS regions at all times (!)". Snapshots are not online.

upvoted 3 times

beginnercloud 1 year, 9 months ago

Selected Answer: C

key words "AWS Regions at all times" so C is correct

upvoted 2 times

fuzzycr 1 year, 10 months ago

Selected Answer: C

key words "AWS Regions at all times"

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: C

By migrating the PostgreSQL database to an RDS for PostgreSQL DB instance and creating a read replica in another AWS Region, you can achieve data availability and online access across multiple Regions. This solution requires less operational overhead compared to managing a PostgreSQL cluster on EC2 instances (Option A) or setting up manual replication using snapshots (Option D). Additionally, Amazon RDS handles the underlying infrastructure and replication setup, reducing the operational complexity for the company.

operational complexity for the company.

Option B, is a valid solution for achieving high availability within a single AWS Region. However, it does not meet the requirement of having the data available and online across multiple AWS Regions at all times, which is specified in the question. The Multi-AZ feature in RDS provides automatic failover within the same Region, but it does not replicate the data to multiple Regions.

upvoted 6 times

mal1903 1 year, 11 months ago

Selected Answer: B

C and D just specify another single region. This does not translate to multiple regions.

B (Multi-AZ) means the solution will be highly available.

The data will be available in multiple regions for both B and C but B is a better solution!

upvoted 1 times

Guru4Cloud 1 year, 8 months ago

its data is available and online across multiple AWS Regions at all times

upvoted 1 times

pentium75 1 year, 5 months ago

"Another single region" plus the original region ARE of course "multiple regions", what else? While B "Multi-AZ" is two AZs in the same region.

upvoted 2 times

MrAWSAssociate 1 year, 11 months ago

Selected Answer: C

Answer B is not right, because "RDS Multi-AZ" always span at least two Availability Zones within a single region and the question requirement RDS DB should be available in multiple regions. Therefore, C is the most suitable answer for this question.

upvoted 3 times

MrAWSAssociate 1 year, 11 months ago

I would like to change my answer to "B". The question has some distractor words: "its data is available and online across multiple AWS Regions at all times". We agree that AWS is a cloud service available online around the world in 99 regions. So the option "B" is the most appropriate answer, since multi-AZ focuses on the availability factor and it has the LEAST amount of operational overhead.

upvoted 1 times

abhishek2021 1 year, 11 months ago

Selected Answer: B

B & C both make data available. However, B is less overhead.

What I think, the question is asking for data availability across multiple regions not for a DR solution. So, RDS being accessible over public IP will do the trick for data being available across regions.

upvoted 1 times

Guru4Cloud 1 year, 8 months ago

Multi-AZ is not the same as Multi-Regional

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #242

Topic 1


A company hosts its web application on AWS using seven Amazon EC2 instances. The company requires that the IP addresses of all healthy EC2 instances be returned in response to DNS queries.

Which policy should be used to meet this requirement?

- A. Simple routing policy
- B. Latency routing policy
- C. Multivalue routing policy **Most Voted**
- D. Geolocation routing policy

Correct Answer: C

Community vote distribution



Comments

cookieMr **Highly Voted** 1 year, 5 months ago

The Multivalue routing policy allows Route 53 to respond to DNS queries with multiple healthy IP addresses for the same resource. This is particularly useful in scenarios where multiple instances are serving the same purpose and need to be load balanced or failover capable. With the Multivalue routing policy, Route 53 returns multiple IP addresses in a random order to distribute the traffic across all healthy instances.

Option A (Simple routing policy) would only return a single IP address in response to DNS queries and does not support returning multiple addresses.

Option B (Latency routing policy) is used to route traffic based on the lowest latency to the resource and does not fulfill the requirement of returning all healthy IP addresses.

Option D (Geolocation routing policy) is used to route traffic based on the geographic location of the user and does not fulfill the requirement of returning all healthy IP addresses.

Therefore, the Multivalue routing policy is the most suitable option for returning the IP addresses of all healthy EC2 instances in response to DNS queries.

upvoted 16 times

LuckyAro **Highly Voted** 1 year, 10 months ago

Selected Answer: C

Use a multivalued answer routing policy to help distribute DNS responses across multiple resources. For example, use multivalued answer routing when you want to associate your routing records with a Route 53 health check. For example, use multivalued answer routing when you need to return multiple values for a DNS query and route traffic to multiple IP addresses.

<https://aws.amazon.com/premiumsupport/knowledge-center/multivalued-versus-simple-policies/>
upvoted 14 times

awsgeek75 Most Recent 11 months, 1 week ago

Selected Answer: C

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html#routing-policy-multivalued>

"Multivalued answer routing policy – Use when you want Route 53 to respond to DNS queries with up to eight healthy records selected at random. You can use multivalued answer routing to create records in a private hosted zone."

Company requires that the IP addresses of "ALL" healthy EC2 instances be returned so C is the only option.
upvoted 5 times

Ruffyt 1 year ago

Use a multivalued answer routing policy to help distribute DNS responses across multiple resources. For example, use multivalued answer routing when you want to associate your routing records with a Route 53 health check. For example, use multivalued answer routing when you need to return multiple values for a DNS query and route traffic to multiple IP addresses.
upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: C

Use Multivalued answer routing policy when you want Route 53 to respond to DNS queries with up to eight healthy records selected at random.
upvoted 2 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: C

C. Multivalued routing policy
upvoted 2 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: A

A. Deploy an AWS Storage Gateway file gateway as a virtual machine (VM) on premises at each clinic
upvoted 1 times

pentium75 11 months, 2 weeks ago

This is for a different question I guess
upvoted 2 times

animefan1 1 year, 5 months ago

multivalued supports health checks
upvoted 2 times

MLCL 1 year, 8 months ago

IP are returned RANDOMLY for multi-value Routing, is this what we want ?
upvoted 4 times

WhericanIstart 1 year, 8 months ago

Selected Answer: C

Multivalued answer routing policy ...answer is C
upvoted 2 times

mhmt4438 1 year, 10 months ago

Selected Answer: C

Answer is C
upvoted 3 times

Aninina 1 year, 10 months ago

Selected Answer: C

Should be C
upvoted 2 times

bamishr 1 year, 11 months ago

Selected Answer: C

<https://www.examtopics.com/discussions/amazon/view/46491-exam-aws-certified-solutions-architect-associate-saa-c02/>
upvoted 2 times

Morinator 1 year, 11 months ago

Selected Answer: C

<https://www.examtopics.com/discussions/amazon/view/46491-exam-aws-certified-solutions-architect-associate-saa-c02/>
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #243

Topic 1

A medical research lab produces data that is related to a new study. The lab wants to make the data available with minimum latency to clinics across the country for their on-premises, file-based applications. The data files are stored in an Amazon S3 bucket that has read-only permissions for each clinic.

What should a solutions architect recommend to meet these requirements?

- A. Deploy an AWS Storage Gateway file gateway as a virtual machine (VM) on premises at each clinic **Most Voted**
- B. Migrate the files to each clinic's on-premises applications by using AWS DataSync for processing.
- C. Deploy an AWS Storage Gateway volume gateway as a virtual machine (VM) on premises at each clinic.
- D. Attach an Amazon Elastic File System (Amazon EFS) file system to each clinic's on-premises servers.

Correct Answer: A*Community vote distribution***Comments****mhmt4438** **Highly Voted** 1 year, 10 months ago**Selected Answer: A**

A. Deploy an AWS Storage Gateway file gateway as a virtual machine (VM) on premises at each clinic

AWS Storage Gateway is a service that connects an on-premises software appliance with cloud-based storage to provide seamless and secure integration between an organization's on-premises IT environment and AWS's storage infrastructure. By deploying a file gateway as a virtual machine on each clinic's premises, the medical research lab can provide low-latency access to the data stored in the S3 bucket while maintaining read-only permissions for each clinic. This solution allows the clinics to access the data files directly from their on-premises file-based applications without the need for data transfer or migration.

upvoted 24 times

cookieMr **Highly Voted** 1 year, 5 months ago

A. It allows the clinics to access the data files stored in the S3 bucket through a file interface. The file gateway caches frequently accessed data locally, reducing latency and providing fast access to the data.

B. It involves transferring the data files from the Amazon S3 bucket to each clinic's on-premises applications using AWS DataSync. While this enables data migration, it may not provide real-time access and may introduce additional latency.

C. It is suitable for block-level access to data rather than file-level access. It may not be the most efficient solution for file-based applications.

D. It involves using Amazon EFS, which is a scalable file storage service, to provide file-level access to the data. However, it may introduce additional complexity and latency compared to using a file gateway solution.

upvoted 8 times

satyaamm Most Recent 3 months, 2 weeks ago

Selected Answer: A

AWS Storage File Gateway is the most suitable here.

upvoted 1 times

pentium75 11 months, 2 weeks ago

Selected Answer: A

A: does exactly that is required here

B: "Migrate", as to MOVE the files out from S3, doesn't make sense

C: Volume Gateway provides iSCSI volumes backed by an object in AWS-managed S3, it does not provide access to S3 objects

D: You can do that but it would have high (not "minimum") latency, and the data is not in that EFS volume, it's in S3

upvoted 5 times

Ruffyit 1 year ago

AWS Storage Gateway is a service that connects an on-premises software appliance with cloud-based storage to provide seamless and secure integration between an organization's on-premises IT environment and AWS's storage infrastructure. By deploying a file gateway as a virtual machine on each clinic's premises, the medical research lab can provide low-latency access to the data stored in the S3 bucket while maintaining read-only permissions for each clinic. This solution allows the clinics to access the data files directly from their on-premises file-based applications without the need for data transfer or migration.

upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: A

The Amazon S3 File Gateway enables you to store and retrieve objects in Amazon Simple Storage Service (S3) using file protocols such as Network File System (NFS) and Server Message Block (SMB). Objects written through S3 File Gateway can be directly accessed in S3.

upvoted 4 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: A

A. Deploy an AWS Storage Gateway file gateway as a virtual machine (VM) on premises at each clinic

upvoted 3 times

Bmarodi 1 year, 6 months ago

Selected Answer: A

Option A meets the requirements.

upvoted 2 times

jaswantn 1 year, 7 months ago

For File-based applications use File Gateway: (Option A)

upvoted 2 times

Grace83 1 year, 8 months ago

Definitely A.

Why are there so many wrong answers by Admins?

upvoted 5 times

maggie135 1 year, 5 months ago

I guess to force us to read and think, so one can't just memorize the answer and go to exam ?)

upvoted 3 times

AlexandraSAA 1 year, 9 months ago

Alessandro333 1 year, 2 months ago

Selected Answer: A

Amazon S3 File Gateway enables you to store file data as objects in Amazon S3 cloud storage for data lakes, backups, and Machine Learning workflows. With Amazon S3 File Gateway, each file is stored as an object in Amazon S3 with a one-to-one mapping between a file and an object.

Volume Gateway provides block storage volumes over iSCSI, backed by Amazon S3, and provides point-in-time backups as Amazon EBS snapshots. Volume Gateway integrates with AWS Backup, an automated and centralized backup service, to protect Storage Gateway volumes.

So it's A

upvoted 5 times

Steve_4542636 1 year, 9 months ago

Selected Answer: A

A for answer

upvoted 2 times

bdp123 1 year, 10 months ago

Selected Answer: A

<https://cloud.in28minutes.com/aws-certification-aws-storage-gateway>

upvoted 2 times

kbaru 1 year, 10 months ago

Selected Answer: A

A. Deploy an AWS Storage Gateway file gateway...

upvoted 2 times

imisioluwa 1 year, 10 months ago

Selected Answer: A

The correct answer is A.

<https://www.knowledgehut.com/tutorials/aws/aws-storage-gateway#:~:text=AWS%20Storage%20Gateway%20helps%20in%20connecting,as%20well%20as%20providing%20data%20security.&text=AWS%20Storage%20Gateway%20helps,as%20providing%20data%20security.&text=Gateway%20helps%20in%20connecting,as%20well%20as%20providing>

<https://docs.aws.amazon.com/storagegateway/latest/vgw/WhatIsStorageGateway.html>

upvoted 2 times

venice1234 1 year, 10 months ago

Selected Answer: C

I think C (Volume Gateway) is correct as it has an option to have Local Storage with Asynchronous sync with S3. This would give low latency access to all local files not just cached/recent files.

upvoted 2 times

pentium75 11 months, 2 weeks ago

Volume gateway provides an iSCSI volume and stores that as a single object in an AWS-managed S3 bucket. It does not provide access to S3 objects.

upvoted 2 times

laicos 1 year, 10 months ago

Selected Answer: A

<https://aws.amazon.com/storagegateway/file/>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #244

Topic 1

A company is using a content management system that runs on a single Amazon EC2 instance. The EC2 instance contains both the web server and the database software. The company must make its website platform highly available and must enable the website to scale to meet user demand.

What should a solutions architect recommend to meet these requirements?

- A. Move the database to Amazon RDS, and enable automatic backups. Manually launch another EC2 instance in the same Availability Zone. Configure an Application Load Balancer in the Availability Zone, and set the two instances as targets.
- B. Migrate the database to an Amazon Aurora instance with a read replica in the same Availability Zone as the existing EC2 instance. Manually launch another EC2 instance in the same Availability Zone. Configure an Application Load Balancer, and set the two EC2 instances as targets.
- C. Move the database to Amazon Aurora with a read replica in another Availability Zone. Create an Amazon Machine Image (AMI) from the EC2 instance. Configure an Application Load Balancer in two Availability Zones. Attach an Auto Scaling group that uses the AMI across two Availability Zones. **Most Voted**
- D. Move the database to a separate EC2 instance, and schedule backups to Amazon S3. Create an Amazon Machine Image (AMI) from the original EC2 instance. Configure an Application Load Balancer in two Availability Zones. Attach an Auto Scaling group that uses the AMI across two Availability Zones.

Correct Answer: C*Community vote distribution*

C (98%)

D

Comments**mhmt4438** **Highly Voted** 1 year, 10 months ago**Selected Answer: C**

C. Move the database to Amazon Aurora with a read replica in another Availability Zone. Create an Amazon Machine Image (AMI) from the EC2 instance. Configure an Application Load Balancer in two Availability Zones. Attach an Auto Scaling group that uses the AMI across two Availability Zones.

This approach will provide both high availability and scalability for the website platform. By moving the database to Amazon Aurora with a read replica in another availability zone, it will provide a failover option for the database. The use of an Application Load Balancer and an Auto Scaling group across two availability zones allows for automatic scaling of the website to meet increased user demand. Additionally, creating an AMI from the original EC2 instance allows for easy replication of the

to meet increased user demand. Additionally, creating an AMI from the original EC2 instance allows for easy replication of the instance in case of failure.

upvoted 23 times

Bmarodi 1 year, 6 months ago

Very good explanations!

upvoted 2 times

cookieMr Highly Voted 1 year, 5 months ago

Selected Answer: C

Option A does not provide a solution for high availability or scalability. Manually launching another EC2 instance in the same AZ may not ensure high availability, as a failure in that AZ would result in downtime.

Option B improves database performance and provides a level of fault tolerance, it does not address the scalability aspect of the website platform.

Option C provides both high availability and fault tolerance. Creating an AMI allows for easy replication of the EC2 instance across AZs. Configuring an ALB in two AZs and attaching an ASG ensures scalability and load distribution across multiple instances.

Option D does not provide the high availability and scalability required by the company. Scheduled backups to S3 address data protection but do not contribute to website availability or scalability.

upvoted 6 times

satyaamm Most Recent 3 months, 2 weeks ago

Selected Answer: C

Aurora with read replica in different region is the most suitable here.

upvoted 1 times

awsgeek75 11 months, 1 week ago

Selected Answer: C

C is the only option via deduction logic based on the assumption the CMS database is Aurora compatible. Other solutions don't promise scaling as much as Aurora solution in option C does.

upvoted 2 times

awsgeek75 10 months, 3 weeks ago

Just to clarify, the question is vague as we don't know anything about the DB types on the CMS so making an assumption that Aurora will work with the CMS.

upvoted 4 times

pentium75 11 months, 2 weeks ago

Selected Answer: C

A and B involve manual steps and do not include scaling (it's just two fixed instances)

D scales the application part but leaves the database on a single EC2 instance which would be neither "highly available" nor "scaleable"

upvoted 5 times

Ruffyt 1 year ago

C. Move the database to Amazon Aurora with a read replica in another Availability Zone. Create an Amazon Machine Image (AMI) from the EC2 instance. Configure an Application Load Balancer in two Availability Zones. Attach an Auto Scaling group that uses the AMI across two Availability Zones.

upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: C

Move the database to Amazon Aurora with a read replica in another Availability Zone. Create an Amazon Machine Image (AMI) from the EC2 instance. Configure an Application Load Balancer in two Availability Zones. Attach an Auto Scaling group that uses the AMI across two Availability Zones.

upvoted 2 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: C

C. Move the database to Amazon Aurora with a read replica in another Availability Zone. Create an Amazon Machine Image (AMI) from the EC2 instance. Configure an Application Load Balancer in two Availability Zones. Attach an Auto Scaling group that uses the AMI across two Availability Zones.

upvoted 2 times

MutiverseAgent 1 year, 4 months ago

Selected Answer: D

The question does not say if the current application is using a relational database, so how we can be sure that it can moved to RDS or aurora as answers A, B & C states? In my opinion the right answer is D.

upvoted 1 times

pentium75 11 months, 2 weeks ago

In D, you "move the database to a separate EC2 instance" BEFORE creating the AMI for the Auto Scaling group. So you'd still have a single EC2 instance running the database, which would meet neither the availability nor the scalability requirement.

upvoted 2 times

animefan1 1 year, 5 months ago

Selected Answer: C

has all options needed for HA

upvoted 2 times

Bmarodi 1 year, 6 months ago

Selected Answer: C

Option C meets the requirements.

upvoted 2 times

ssoffline 1 year, 6 months ago

Why not D?

Are we just assuming that there will be no write to the db?

upvoted 1 times

antropaws 1 year, 6 months ago

Selected Answer: C

Absolutely C.

upvoted 2 times

Aninina 1 year, 10 months ago

Selected Answer: C

C: This will allow the website platform to be highly available by using Aurora, which provides automatic failover and replication. Additionally, by creating an AMI from the original EC2 instance, the Auto Scaling group can automatically launch new instances in multiple availability zones and use the Application Load Balancer to distribute traffic across them. This way, the website will be able to handle the increased traffic, and will be less likely to go down due to a single point of failure.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #245

Topic 1

A company is launching an application on AWS. The application uses an Application Load Balancer (ALB) to direct traffic to at least two Amazon EC2 instances in a single target group. The instances are in an Auto Scaling group for each environment. The company requires a development environment and a production environment. The production environment will have periods of high traffic.


Which solution will configure the development environment MOST cost-effectively?

- A. Reconfigure the target group in the development environment to have only one EC2 instance as a target.
- B. Change the ALB balancing algorithm to least outstanding requests.
- C. Reduce the size of the EC2 instances in both environments.

D. Reduce the maximum number of EC2 instances in the development environment's Auto Scaling group. **Most Voted**

Correct Answer: D

Community vote distribution

**Comments**

mhmt4438 **Highly Voted** 2 years, 4 months ago

Selected Answer: D

D. Reduce the maximum number of EC2 instances in the development environment's Auto Scaling group

This option will configure the development environment in the most cost-effective way as it reduces the number of instances running in the development environment and therefore reduces the cost of running the application. The development environment typically requires less resources than the production environment, and it is unlikely that the development environment will have periods of high traffic that would require a large number of instances. By reducing the maximum number of instances in the development environment's Auto Scaling group, the company can save on costs while still maintaining a functional development environment.

upvoted 20 times

JayBee65 2 years, 4 months ago

No, it will not reduce the number of instances being used, since a minimum of 2 will be used at all times.

upvoted 10 times

pentium75 1 year, 5 months ago

But it will keep the number of instances at two, while the production environment has "AT LEAST two".

upvoted 4 times

pentium75 **Highly Voted** 1 year, 5 months ago

Selected Answer: D

A modifies only the ALB target group (= directs traffic only to one node), but does not affect the number of nodes (and the cost)

B balances load between nodes but does not affect the cost

C impacts the prod environment so that would be unable to handle its "periods of high traffic"

D makes sure that the dev environment will not scale to more than 2 instances, as does the prod environment

upvoted 10 times

Yak_Yeti **Most Recent** 1 month, 3 weeks ago

Selected Answer: A

you cannot get cheaper than one, and D says to reduce the maximum and does not mention the minimum.

upvoted 1 times

jerryI 2 months, 2 weeks ago

Selected Answer: D

analysis/explanation by copilot

> A: you reduce the target group instance, doesn't mean the other one is shut down / terminated (sound tricky but could be true, it didn't mention the other machine is terminated)

> D: that make sure only 1 machine is always up

in that case then D would make more sense (disregard the high availability concern, A is still more costly then)

upvoted 1 times

SirDNS 2 months, 2 weeks ago

Selected Answer: A

A. This is the most cost-effective approach for a development environment. Since development environments typically don't need the same level of high availability or performance as production, reducing the number of target EC2 instances to one minimizes costs.

D. Reducing the maximum size of the Auto Scaling group can limit scaling, but it doesn't guarantee cost savings if those instances are running. Reconfiguring the target group to use only one instance is a more direct way to reduce costs in the development environment.

upvoted 1 times

zdi561 4 months ago

Selected Answer: A

D is not right, if you decrease max number it will not meet the high traffic requirement. Actually it does not matter because the ASG will automatically adjust the running target. A really means set the max, min and desired to one which can save a little bit money, The requirement for two really means the min is 2 in prod. max could be higher than 2.

upvoted 1 times

XXXXXINN 8 months, 3 weeks ago

chatGPT sucks too. it says A, but A obviously just says reduced the number of Target Group, reduce Target Group does not mean the reduce of EC2 instances themselves, so there is no cost saved at all. Thus DDDDDDDDDDD....

upvoted 2 times

LeonSauveterre 6 months, 2 weeks ago

Weird... ChatGPT says to choose D when I asked about this question. It said option A compromises high availability because ALBs require at least two targets for redundancy and fault tolerance, so not a valid option due to the risk of downtime if the single instance becomes unavailable, while option D provides targeted cost optimization for the development environment without compromising availability or production performance.

upvoted 1 times

babayomi 7 months, 1 week ago

Answer is A, option A says reduce(rec9nfigure) the instance in target group to one, makes sense. If a group has only implemented instance, it's not autoscaling.

upvoted 1 times

MatAlves 9 months ago

Selected Answer: D

So, in short, the question asks for a way to reduce cost wasted with the dev env, since it's resources are being underused.

(A) Target group vs (B) Auto Scaling Group.

Reducing the target group won't affect the number of "nodes" (instances), cost will stay the same. To eliminate the excess of ec2 instances in the dev env, you actually need to reduce the auto scaling group.

upvoted 3 times

TheFivePips 1 year, 3 months ago

Selected Answer: D

The application uses an Application Load Balancer (ALB) to direct traffic to at least two Amazon EC2 instances in a single target group

You are required to keep at least two instances in each target group. A sets it to one, which would be more cost effective, but doesn't meet the requirement.

upvoted 4 times

Priyapani 1 year, 4 months ago

Selected Answer: D

In the question it is said minimum it should have 2 instances in Target group. So in development group we can reduce the the target group.

In option A. It is said it will have only one instance in development group that doesn't match to our question

upvoted 5 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

B and C don't actually save any cost without impacting performance during high traffic on production.

A and D are basically same thing but A enforces a limit of one EC2 instance which is not acceptable as the question asks:

"Application Load Balancer (ALB) to direct traffic to at least two Amazon EC2 instances in a single target group. The instances are in an Auto Scaling group for each environment"

Hence D is the only valid answer.

upvoted 6 times

ddaannndann 1 year, 5 months ago

The most cost-effective solution is to reconfigure the target group in the development environment to have only one EC2 instance as a target. This will ensure that the development environment only uses the resources that it needs, which will save the company money.

The other solutions are not as cost-effective. Changing the ALB balancing algorithm to least outstanding requests will not reduce the number of EC2 instances that are used, and it may actually increase the amount of traffic that is directed to each instance. Reducing the size of the EC2 instances will also not reduce the number of instances that are used, and it may actually make the application slower. Reducing the maximum number of EC2 instances in the development environment's Auto Scaling group will only reduce the number of instances that are used when the traffic is high, and it will not reduce the number of instances that are used on average.

Therefore, the most cost-effective solution is to reconfigure the target group in the development environment to have only one EC2 instance as a target.

upvoted 1 times

pentium75 1 year, 5 months ago

No, you're confusing "target group" (of the ALB) with "Auto Scaling group". Answer A will direct ALB traffic only to one node, but it does not affect the number of nodes in any way (it will still be "at least two").

upvoted 4 times

chasingsummer 1 year, 5 months ago

Selected Answer: A

By reconfiguring the target group in the development environment to have only one EC2 instance as a target, it reduces the number of instances handling the development environment's traffic. This ensures the minimum setup required for the development environment's functionality without incurring unnecessary costs associated with multiple instances.

This solution optimizes costs by scaling down the infrastructure specifically in the development environment where lower traffic or fewer resources might be acceptable for testing or development purposes, thus reducing unnecessary expenses related to running multiple instances.

upvoted 1 times

pentium75 1 year, 5 months ago

No, you're confusing "target group" (of the ALB) with "Auto Scaling group". Answer A will direct ALB traffic only to one node, but it does not affect the number of nodes in any way (it will still be "at least two").

upvoted 7 times

MiniYang 1 year, 5 months ago

Selected Answer: D

the correct answer is D. This is from Amazon Q : The most cost-effective way to configure the development environment would be to reduce the maximum number of EC2 instances in the development environment's Auto Scaling group (Option D). The most cost-effective way to configure the development environment would be to reduce the maximum number of EC2 instances in the development environment's Auto Scaling group (Option D).The most cost-effective way to configure the development environment would be to reduce the maximum number of EC2 instances in the development environment's Auto Scaling group (Option D).

upvoted 4 times

meowruki 1 year, 6 months ago

Selected Answer: A

Option D: Reducing the maximum number of EC2 instances in the development environment's Auto Scaling group could limit scalability but might not directly optimize costs. Min can still be the same number of EC2

upvoted 1 times

saymolet 1 year, 3 months ago

"might not directly optimize." No, it does, you're paying less for fewer machines. This is the most direct cost optimisation practice there could ever be. The correct answer is D

upvoted 1 times

pentium75 1 year, 5 months ago

No, you're confusing "target group" (of the ALB) with "Auto Scaling group". Answer A will direct ALB traffic only to one node, but it does not affect the number of nodes in any way (it will still be "at least two").

upvoted 4 times

Chef_couincouin 1 year, 6 months ago

Answer is A but I'm not agree. We use only one instance with A and D.
But with D, by default, instance is terminated whereas with A, instance still exist.
Answer should be D

upvoted 3 times

ravinperera 1 year, 7 months ago

Selected Answer: D

This option is specific to the development environment and focuses on reducing the number of instances that can be spun up during scaling events. This means cost savings because fewer instances will be used even if the scaling policies are triggered. Given the goal to configure the development environment in the most cost-effective way, without compromising the production environment, the best option is D

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #246

Topic 1

A company runs a web application on Amazon EC2 instances in multiple Availability Zones. The EC2 instances are in private subnets. A solutions architect implements an internet-facing Application Load Balancer (ALB) and specifies the EC2 instances as the target group. However, the internet traffic is not reaching the EC2 instances.

How should the solutions architect reconfigure the architecture to resolve this issue?

- A. Replace the ALB with a Network Load Balancer. Configure a NAT gateway in a public subnet to allow internet traffic.
- B. Move the EC2 instances to public subnets. Add a rule to the EC2 instances' security groups to allow outbound traffic to 0.0.0.0/0.
- C. Update the route tables for the EC2 instances' subnets to send 0.0.0.0/0 traffic through the internet gateway route. Add a rule to the EC2 instances' security groups to allow outbound traffic to 0.0.0.0/0.
- D. Create public subnets in each Availability Zone. Associate the public subnets with the ALB. Update the route tables for the public subnets with a route to the private subnets. **Most Voted**

Correct Answer: D*Community vote distribution*

D (86%)

Other (14%)

Comments**ktulu2602** **Highly Voted** 1 year, 9 months ago

I think either the question or the answers are not formulated correctly because of this document:

<https://docs.aws.amazon.com/prescriptive-guidance/latest/load-balancer-stickiness/subnets-routing.html>

A - Might be possible but it's quite impractical

B - Not needed as the setup described should work as is provided the SGs of the EC2 instances accept traffic from the ALB

C - Update the route tables for the EC2 instances' subnets to send 0.0.0.0/0 traffic through the internet gateway route - not needed as the EC2 instances would receive the traffic from the ALB ENIs. Add a rule to the EC2 instances' security groups to allow outbound traffic to 0.0.0.0/0 - the default behaviour of the SG is to allow outbound traffic only.

D - Create public subnets in each Availability Zone. Associate the public subnets with the ALB - if it's a internet facing ALB these should already be in place. Update the route tables for the public subnets with a route to the private subnets - no need as the local prefix entry in the route tables would take care of this point

I'm 110% sure the question or answers or both are wrong. Prove me wrong! :)

upvoted 21 times

ktulu2602 1 year, 9 months ago

unluckyducky 1 year, 8 months ago

Completely agreed, I was looking for an option to allow HTTPS traffic on port 443 from the ALB to the EC2 instance's security group.

Either the question or the answers are wrong.
upvoted 8 times

bdp123 **Highly Voted** 1 year, 9 months ago

Selected Answer: D

I change my answer to 'D' because of following link:
<https://aws.amazon.com/premiumsupport/knowledge-center/public-load-balancer-private-ec2/>
upvoted 15 times

JackyCCK **Most Recent** 8 months, 1 week ago

Can a EC2 in the private subnet sends traffic to the internet through ELB without using NAT gateway/instance?
if it's only about responses to requests coming through the ELB then no, you don't need NAT. If on the other hand you expect that your instances will need to initiate connections as oppose to just sending responses to the requests then yes, you will have to use NAT. The key point is that response to a request is not a new connection and will be sent to where it came from, i.e. to the ELB.

<https://serverfault.com/questions/986447/can-a-ec2-in-the-private-subnet-sends-traffic-to-the-internet-through-elb-withou>
upvoted 3 times

sidharthwader 9 months ago

D looks the best but still it must have a internet gateway and once it has internet gateway we must add the route table for private subnet to talk to the public subnet so by using the it should be able to access. I don't think lb can act like internet gateway
upvoted 2 times

bujuman 9 months, 2 weeks ago

Selected Answer: D

Considering these statements:
-The EC2 instances are in private subnets.
- However, the internet traffic is not reaching the EC2 instances.
A reliable solution is D according to following link:
<https://repost.aws/knowledge-center/public-load-balancer-private-ec2>
Answer C could not satisfy the requirements because only outbound traffic rules are mentionned
upvoted 2 times

pentium75 11 months, 2 weeks ago

Selected Answer: D

A - "NAT gateway" is "to allow [outbound] internet traffic", but this is about inbound traffic
B - This is about outbound traffic while the problem is inbound
C - This is about outbound traffic while the problem is inbound
D - Sounds correct, though the "update the route tables" should not be required if both subnets are in same VPC
upvoted 7 times

awsgeek75 10 months, 3 weeks ago

D is the "least wrong" answer here. I was also confused by the route table part and thought I was missing something critical in the question.
upvoted 4 times

David_Ang 1 year, 1 month ago

Selected Answer: A

this is a bad formulated question with gaps, but my reason tells me that if you want to connect something from a private subnet to internet you need a NAT (instance or gateway, bastion).
Creating public subnets in each Availability Zone and associating them with the Application Load Balancer (ALB) won't resolve the problem of allowing internet traffic to reach the private EC2 instances. Public subnets are typically used when you want your EC2 instances to have direct internet access, not when you want to keep them in private subnets with indirect access through a load balancer.
upvoted 3 times

vijaykamal 1 year, 2 months ago

Selected Answer: D

ption A (replace ALB with Network Load Balancer and add a NAT gateway) is not the most straightforward solution because it changes the load balancer type and introduces a NAT gateway, which might be unnecessary if the goal is to use an ALB for web traffic. ALBs are commonly used for internet-facing web applications.

Option B (move EC2 instances to public subnets and modify security group rules) involves placing instances in public subnets, which is generally not recommended for security reasons. Additionally, it suggests modifying security group rules for outbound traffic, which might not be the best practice to resolve the issue.

Option C (update route tables and security group rules) addresses the route table update, but it also suggests moving instances to public subnets, which is not ideal from a security perspective.

upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: D

Create public subnets in each Availability Zone. Associate the public subnets with the ALB. Update the route tables for the public subnets with a route to the private subnets.

upvoted 3 times

Its_SaKar 1 year, 2 months ago

Selected Answer: D

Option A is incorrect Internet traffic is http and https so it cant be configured to NLB
Option B and option C is incorrect because sening 0.0.0.0/0 is not best practices

Option D is correct because its the only option left. and updating the route tables for the public subnets with a route to the private subnets ensures internet access to EC2 instances in private subnet.

upvoted 3 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: D

D. is the correct solution. By creating public subnets and associating them with the ALB, inbound internet traffic can reach the ALB. The route tables for the public subnets are updated to include a route to the private subnets, allowing traffic to reach the EC2 instances in the private subnets. This setup enables secure access to the application while allowing internet traffic to reach the EC2 instances through the ALB.

upvoted 4 times

A1975 1 year, 4 months ago

Selected Answer: D

<https://docs.aws.amazon.com/vpc/latest/userguide/vpc-example-private-subnets-nat.html>

upvoted 3 times

cookieMr 1 year, 5 months ago

Selected Answer: D

A. suggests using a different type of load balancer and configuring a NAT gateway, but it does not address the issue of internet traffic reaching the EC2 instances.

B. suggests exposing the EC2 instances to the public internet, which may pose security risks and does not address the issue of inbound internet traffic reaching the instances.

C. suggests configuring the EC2 instances to have outbound internet access, but it does not solve the problem of inbound internet traffic reaching the instances.

D. is the correct solution. By creating public subnets and associating them with the ALB, inbound internet traffic can reach the ALB. The route tables for the public subnets are updated to include a route to the private subnets, allowing traffic to reach the EC2 instances in the private subnets. This setup enables secure access to the application while allowing internet traffic to reach the EC2 instances through the ALB.

upvoted 6 times

Vinhkewl 1 year, 5 months ago

Should be C

It would normally make sense to segregate your ALBs into public or private zones by security group and target group, but this

is configuration rather than architectural placement - there is nothing preventing you from adding a rule to route specific paths or ports to a public subnet from an ALB that has until then been serving private subnets only.

upvoted 1 times

pentium75 11 months, 2 weeks ago

C allows the EC2 instances to be accessed directly from the Internet, which we don't want. It's the ALB (not the Internet) that can't access them. We must make sure that the ALB can be reached from the Internet and that the EC2 instances can be reached from the ALB.

upvoted 2 times

Abrar2022 1 year, 6 months ago

Selected Answer: D

To attach Amazon EC2 instances that are located in a private subnet, first create public subnets

upvoted 5 times

Bmarodi 1 year, 6 months ago

Selected Answer: D

I vote with the option D.

upvoted 2 times

antropaws 1 year, 6 months ago

D is not quite accurate because subnets in a VPC have a local route by default, meaning that all subnets are able to communicate with each other: "Every route table contains a local route for communication within the VPC. This route is added by default to all route tables". This question is poorly formulated.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #247

Topic 1


A company has deployed a database in Amazon RDS for MySQL. Due to increased transactions, the database support team is reporting slow reads against the DB instance and recommends adding a read replica.

Which combination of actions should a solutions architect take before implementing this change? (Choose two.)

- A. Enable binlog replication on the RDS primary node.
- B. Choose a failover priority for the source DB instance.
- C. Allow long-running transactions to complete on the source DB instance. **Most Voted**
- D. Create a global table and specify the AWS Regions where the table will be available.
- E. Enable automatic backups on the source instance by setting the backup retention period to a value other than 0. **Most Voted**

Correct Answer: CE

Community vote distribution



Comments

fkie4 **Highly Voted** 2 years, 3 months ago

Who would know this stuff man...
upvoted 129 times

presetacsing 2 years ago

exactly
upvoted 2 times

MNotABot 1 year, 11 months ago

"Allow long-running transactions to complete on the source DB instance." --. Makes sense / Also a backup before changing anything again made a sense.
upvoted 3 times

foha2012 1 year, 4 months ago

Just take an intelligent guess. Eliminate 2 wrong answers and you will have a 50percent success chance.

upvoted 3 times

KelvinEM Highly Voted 2 years, 4 months ago

C,E

"An active, long-running transaction can slow the process of creating the read replica. We recommend that you wait for long-running transactions to complete before creating a read replica. If you create multiple read replicas in parallel from the same source DB instance, Amazon RDS takes only one snapshot at the start of the first create action.

When creating a read replica, there are a few things to consider. First, you must enable automatic backups on the source DB instance by setting the backup retention period to a value other than 0. This requirement also applies to a read replica that is the source DB instance for another read replica"

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_ReadRepl.html

upvoted 52 times

c12ab95 Most Recent 2 weeks, 6 days ago

Selected Answer: CE

C. Allow long-running transactions to complete on the source DB instance.

Active long-running transactions can delay read replica creation. AWS recommends resolving these first to avoid replication delays or failures .

E. Enable automatic backups on the source instance by setting the backup retention period to a value other than 0.

Automatic backups are mandatory for read replica creation. Enabling backups activates binary logging (binlog) on the source, which is required for replication

upvoted 1 times

Yak_Yeti 1 month, 3 weeks ago

Selected Answer: AE

To create a read replica for Amazon RDS for MySQL, the following prerequisites must be met:

A. Enable binlog replication on the RDS primary node

Required. RDS read replicas rely on MySQL binary log (binlog) replication.

This setting is not enabled by default. You must set the DB parameter binlog_format to ROW and enable binary logging by modifying the parameter group.

E. Enable automatic backups

Required. Enabling automatic backups by setting backup retention > 0 is necessary to support point-in-time recovery and to allow the creation of read replicas.

RDS uses the backup snapshot as a baseline to create the initial read replica.

Why not "C. Allow long-running transactions to complete"?

Not a prerequisite for creating a read replica.

Might improve performance temporarily, but it's unrelated to read replica setup.

upvoted 2 times

Yak_Yeti 1 month, 3 weeks ago

Moderator - please delete

Enabling binary log replication on the RDS primary node is indeed automatically handled by Amazon RDS when you create a read replica for MySQL or MariaDB. You do not need to manually configure this beforehand.

upvoted 1 times

Rcosmos 4 months, 3 weeks ago

Selected Answer: UE

Opção A: Habilite a replicação de log binário no nó primário do RDS

A replicação de log binário é necessária para habilitar réplicas de leitura no RDS for MySQL. O banco de dados de origem precisa capturar e armazenar logs binários que serão usados para replicar alterações para as réplicas de leitura,

E: Habilite backups automáticos na instância de origem definindo o período de retenção de backup como um valor diferente de 0 Backups automáticos devem estar habilitados na instância de origem para criar réplicas de leitura. Isso garante que a réplica possa ser criada a partir de um ponto consistente do banco de dados de origem

upvoted 1 times

derekxxxxxxx 6 months ago

Selected Answer: CE

Allow long-running transactions to complete on the source DB instance.

upvoted 1 times

FlyingHawk 6 months, 2 weeks ago

Selected Answer: CE

Based on the official documentation, the correct steps to prepare for creating a read replica are:

- A. Enable binlog replication on the RDS primary node.
 - C. Allow long-running transactions to complete on the source DB instance.
 - E. Enable automatic backups on the source instance by setting the backup retention period to a value other than 0.
- These ensure smooth and efficient setup while adhering to AWS best practices.

For more detailed information, you can refer to the AWS documentation:

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_ReadRepl.Create.html

A (Enable binlog replication) is not a required manual step in Amazon RDS.

Amazon RDS automatically manages this requirement during the creation of the read replica for MySQL and MariaDB.

upvoted 2 times

LeonSauveterre 6 months, 2 weeks ago

Selected Answer: CE

Answer: CE.

A: In Amazon RDS, binary logging is automatically enabled when creating a read replica. You don't need to enable it manually.

B: Read replicas are used for improving read performance, not for failover purposes.

C: This ensures a clean state before replication begins.

D: Global tables are part of Amazon DynamoDB, not Amazon RDS.

E: RDS uses the automated backup feature to take a snapshot of the primary DB instance and initialize the replica.

upvoted 2 times

tonybuivannghia 8 months, 1 week ago

Selected Answer: CE

A isn't correct because binary log is just for external DB instance.

upvoted 3 times

ccceb01 9 months, 3 weeks ago

Answer is C and E

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_ReadRepl.html

upvoted 2 times

1166ae3 11 months, 2 weeks ago

Selected Answer: AE

A. Enable binlog replication on the RDS primary node:

Direct Impact: Enabling binlog replication is crucial for setting up read replicas, which will directly help in distributing the read load and improving read performance.

E. Enable automatic backups on the source instance by setting the backup retention period to a value other than 0:

Direct Impact: Automatic backups are necessary to ensure data integrity when creating read replicas. This setup is critical for maintaining consistent and reliable replicas.

upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: CE

B and D don't have anything to do with the question.

E is a must have before doing major architecture changes

A is not something you need to do explicitly when creating read replicas as it is managed by RDS

C makes sense

* I think the options are really badly worded which makes it confusing. I doubt this is a real question.

upvoted 4 times

awsgeek75 1 year, 5 months ago

Also, to add, binlog replication is needed if you are replicating to a non RDS instance. This is why I think the question is badly phrased as it does not specify the location of read replica.

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/MySQL.Procedural.Importing.External.Repl.html>

upvoted 1 times

Ruffyt 1 year, 6 months ago

An active, long-running transaction can slow the process of creating the read replica. We recommend that you wait for long-running transactions to complete before creating a read replica. If you create multiple read replicas in parallel from the same source DB instance, Amazon RDS takes only one snapshot at the start of the first create action.

When creating a read replica, there are a few things to consider. First, you must enable automatic backups on the source DB instance by setting the backup retention period to a value other than 0. This requirement also applies to a read replica that is the source DB instance for another read replica

upvoted 2 times

meowruki 1 year, 6 months ago

Selected Answer: AC

To improve the read performance of a database in Amazon RDS for MySQL by adding a read replica, you should take the following actions:

Enable binlog replication on the RDS primary node: This allows the primary node to stream its binary logs to the read replica, enabling data replication.

A. Enable binlog replication on the RDS primary node.

Allow long-running transactions to complete on the source DB instance: Before creating a read replica, it's advisable to let any long-running transactions complete to ensure consistency between the source and the replica.

C. Allow long-running transactions to complete on the source DB instance.

The other options are not directly related to setting up a read replica:

upvoted 1 times

meowruki 1 year, 6 months ago

B. Choose a failover priority for the source DB instance: Failover priority is more relevant in a Multi-AZ setup where automatic failover might occur. It's not directly related to creating a read replica.

D. Create a global table and specify the AWS Regions where the table will be available: Global tables are used for cross-region replication, but they are not directly related to setting up a read replica for improved read performance.

E. Enable automatic backups on the source instance by setting the backup retention period to a value other than 0: While it's a good practice to have backups enabled, it is not a prerequisite for creating a read replica.

Therefore, the most appropriate actions are A and C.

upvoted 1 times

xdkonorek2 1 year, 6 months ago

Selected Answer: AE

A - it's essential for continuous replication

E - it's essential for setting up replication, initial data in replica is based on latest backup

other options:

B - we're not designing for HA, and it's related to multi-AZ RDS deployments

C - is this needed for adding read replica?

D - it's not a dynamodb to create global table

upvoted 5 times

vijaykamal 1 year, 8 months ago

Selected Answer: CE

A. Enabling binlog replication is not something you need to do manually before creating a read replica. Amazon RDS for MySQL manages replication internally, and it's not necessary to enable binlog replication explicitly.

B. Choosing a failover priority is related to Multi-AZ configurations and automatic failover, but it is not specifically required when adding a read replica.

D. Creating a global table and specifying AWS Regions is related to Aurora Global Databases, which is not the same as creating a read replica for a standard RDS instance.

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: CE

**C. Long-running transactions can prevent the read replica from catching up with the source DB instance. Allowing these transactions to complete before creating the read replica can help ensure that the replica is able to stay synchronized with the source.

**E. Automatic backups must be enabled on the source DB instance for read replicas to be created. This is done by setting the backup retention period to a value other than 0.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #248

Topic 1

A company runs analytics software on Amazon EC2 instances. The software accepts job requests from users to process data that has been uploaded to Amazon S3. Users report that some submitted data is not being processed Amazon CloudWatch reveals that the EC2 instances have a consistent CPU utilization at or near 100%. The company wants to improve system performance and scale the system based on user load.

What should a solutions architect do to meet these requirements?

- A. Create a copy of the instance. Place all instances behind an Application Load Balancer.
- B. Create an S3 VPC endpoint for Amazon S3. Update the software to reference the endpoint.
- C. Stop the EC2 instances. Modify the instance type to one with a more powerful CPU and more memory. Restart the instances.
- D. Route incoming requests to Amazon Simple Queue Service (Amazon SQS). Configure an EC2 Auto Scaling group based on queue size. Update the software to read from the queue. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

mhmt4438 **Highly Voted** 1 year, 10 months ago

Selected Answer: D

D. Route incoming requests to Amazon Simple Queue Service (Amazon SQS). Configure an EC2 Auto Scaling group based on queue size. Update the software to read from the queue.

By routing incoming requests to Amazon SQS, the company can decouple the job requests from the processing instances. This allows them to scale the number of instances based on the size of the queue, providing more resources when needed. Additionally, using an Auto Scaling group based on the queue size will automatically scale the number of instances up or down depending on the workload. Updating the software to read from the queue will allow it to process the job requests in a more efficient manner, improving the performance of the system.

upvoted 13 times

b1e2cbe **Most Recent** 7 months, 3 weeks ago

D is correct

upvoted 2 times

Pangian 9 months, 4 weeks ago

Selected Answer: D

The based on the queue size doesn't seem a perfect approach though

upvoted 2 times

awsgeek75 11 months, 1 week ago

Selected Answer: D

D: Because this whole exam seems to be selling more and more SQS solutions...

upvoted 3 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: D

Route incoming requests to Amazon Simple Queue Service (Amazon SQS). Configure an EC2 Auto Scaling group based on queue size. Update the software to read from the queue

upvoted 2 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: D

D. Route incoming requests to Amazon Simple Queue Service (Amazon SQS). Configure an EC2 Auto Scaling group based on queue size. Update the software to read from the queue.

upvoted 2 times

Kill3rasp3r 1 year, 3 months ago

Selected Answer: D

I would vote A if it was ALB targeting an EC2 auto scaling group.

I would vote D if the auto scaling group was based on CPU utilization rather than queue size.

So I think both answers are wrong but D is okay enough.

upvoted 4 times

cookieMr 1 year, 5 months ago

Selected Answer: D

A. Creating a copy of the instance and placing all instances behind an ALB does not address the high CPU utilization issue or provide scalability based on user load.

B. Creating an S3 VPC endpoint for S3 and updating the software to reference the endpoint improves network performance but does not address the high CPU utilization or provide scalability based on user load.

C. Stopping the EC2 instances and modifying the instance type to one with a more powerful CPU and more memory may improve performance, but it does not address scalability based on user load.

D. Routing incoming requests to SQS, configuring an EC2 ASG based on queue size, and updating the software to read from the queue improves system performance and provides scalability based on user load.

Therefore, option D is the correct choice as it addresses the high CPU utilization, improves system performance, and enables scalability based on user load.

upvoted 3 times

WherecanIstart 1 year, 8 months ago

Selected Answer: D

Autoscaling Group and SQS solves the problem.

SQS - Decouples the process

ASG - Autoscales the EC2 instances based on usage

upvoted 2 times

ak1ak 1 year, 9 months ago

Selected Answer: A

its definitely A

upvoted 1 times

wRhIH 1 year, 6 months ago

You don't "scale the system by load" by choosing A
upvoted 4 times

AHUI 1 year, 10 months ago

D is correct. Decouple the process. autoscale the EC2 based on query size. best choice
upvoted 4 times

Aninina 1 year, 10 months ago

I think it's A " A. Create a copy of the instance. Place all instances behind an Application Load Balancer.
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #249

Topic 1

A company is implementing a shared storage solution for a media application that is hosted in the AWS Cloud. The company needs the ability to use SMB clients to access data. The solution must be fully managed.

Which AWS solution meets these requirements?

- A. Create an AWS Storage Gateway volume gateway. Create a file share that uses the required client protocol. Connect the application server to the file share.
- B. Create an AWS Storage Gateway tape gateway. Configure tapes to use Amazon S3. Connect the application server to the tape gateway.
- C. Create an Amazon EC2 Windows instance. Install and configure a Windows file share role on the instance. Connect the application server to the file share.
- D. Create an Amazon FSx for Windows File Server file system. Attach the file system to the origin server. Connect the application server to the file system. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

Morinator **Highly Voted** 2 years, 4 months ago

Selected Answer: D

SMB + fully managed = fsx for windows imo
upvoted 20 times

devonwho **Highly Voted** 2 years, 4 months ago

Selected Answer: D

Amazon FSx has native support for Windows file system features and for the industry-standard Server Message Block (SMB) protocol to access file storage over a network.

<https://docs.aws.amazon.com/fsx/latest/WindowsGuide/what-is.html>
upvoted 7 times

• **devonwho** 2 years, 4 months ago

Leonsauveterre 6 months, 2 weeks ago

As implied in the question, "The solution must be fully managed". Storage Gateway's SMB support is valid but not the best choice for this fully AWS-hosted and fully managed requirement.

Storage Gateway requires the deployment of a gateway appliance (hardware or virtual). While AWS manages the back end (S3), setting up and maintaining the gateway appliance requires some operational effort, so it's not fully managed compared to FSx for Windows File Server.

Storage Gateway is ideal for hybrid environments but adds unnecessary complexity if the workload is fully cloud-based. Not to mention that Storage Gateway File Gateway is focused on storing objects in S3, not providing a shared file system for direct collaboration or use by applications.

upvoted 2 times

kg508 Most Recent 3 months ago

Selected Answer: D

□□□□ windows□ □□ □□□□ □ □□ fsx for windows□□ □□ □□□
D□ □□□ □□□□□ □□ □□□□□□

upvoted 1 times

0xE8D4A51000 11 months, 2 weeks ago

Selected Answer: A

All the answers are wrong here by people. We don't know whether the user is using Windows based applications. AWS Storage Gateway also supports SMB protocol. This is the answer.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: D

A: Volume Gateway provides virtual disks iSCSI, not SMB
B: Tape Gateway provides virtual tapes via iSCSI, not SMB
C: Not "fully managed"

upvoted 5 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: D

SMB = Amazon FSx for Windows File Server

upvoted 5 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

D. Create an Amazon FSx for Windows File Server file system. Attach the file system to the origin server. Connect the application server to the file system

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

All who selected D. are correct - see more details from our community

upvoted 2 times

animefan1 1 year, 11 months ago

Selected Answer: D

Fsx is fully managed. Plus it supports SMB protocol

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: D

A. involves using Storage Gateway, but it does not specifically mention support for SMB clients. It may not meet the requirement of using SMB clients to access data.

B. involves using Storage Gateway with tape gateway configuration, which is primarily used for archiving data to S3. It does not provide native support for SMB clients to access data.

C. involves manually setting up and configuring a Windows file share on an EC2 Windows instance. While it allows SMB clients to access data, it is not a fully managed solution as it requires manual setup and maintenance.

D. involves creating an FSx for Windows File Server file system, which is a fully managed Windows file system that supports SMB clients. It provides an easy-to-use shared storage solution with native SMB support.

Based on the requirements of using SMB clients and needing a fully managed solution, option D is the most suitable choice.

upvoted 6 times

LuckyAro 2 years, 4 months ago

Selected Answer: D

Amazon FSx for Windows File Server file system

upvoted 2 times

techhb 2 years, 4 months ago

amazon fsx for smb connectivity.

upvoted 2 times

Aninina 2 years, 4 months ago

Selected Answer: D

FSX is the ans

upvoted 2 times

mhmt4438 2 years, 4 months ago

Selected Answer: D

<https://www.examttopics.com/discussions/amazon/view/81115-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 3 times

bamishr 2 years, 4 months ago

Selected Answer: D

D. Create an Amazon FSx for Windows File Server file system. Attach the file system to the origin server. Connect the application server to the file system.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #250

Topic 1


A company's security team requests that network traffic be captured in VPC Flow Logs. The logs will be frequently accessed for 90 days and then accessed intermittently.

What should a solutions architect do to meet these requirements when configuring the logs?

- A. Use Amazon CloudWatch as the target. Set the CloudWatch log group with an expiration of 90 days
- B. Use Amazon Kinesis as the target. Configure the Kinesis stream to always retain the logs for 90 days.
- C. Use AWS CloudTrail as the target. Configure CloudTrail to save to an Amazon S3 bucket, and enable S3 Intelligent-Tiering.
- D. Use Amazon S3 as the target. Enable an S3 Lifecycle policy to transition the logs to S3 Standard-Infrequent Access (S3 Standard-IA) after 90 days. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

cookieMr **Highly Voted** 1 year, 11 months ago

Selected Answer: D

A. suggests using CloudWatch as the target for VPC Flow Logs. However, it does not provide a mechanism for managing the retention of the logs for 90 days and then accessing them intermittently.

B. suggests using Kinesis as the target for VPC Flow Logs. While it can retain the logs for 90 days, it does not address the requirement for intermittent access to the logs.

C. suggests using CloudTrail as the target for VPC Flow Logs. However, CloudTrail is designed for auditing and monitoring API activity, not for capturing network traffic logs. It does not meet the requirement of capturing VPC Flow Logs.

D. suggests using S3 as the target for VPC Flow Logs and leveraging S3 Lifecycle policies to transition the logs to a cost-effective storage class after 90 days. It meets the requirement of retaining the logs for 90 days and provides the flexibility for intermittent access while optimizing storage costs.

upvoted 10 times

LuckyAro **Highly Voted** 2 years, 4 months ago

Selected Answer: D

D is the correct answer.

upvoted 6 times

manabpokhrel7 Most Recent 7 months, 3 weeks ago

Selected Answer: D

D is the correct answer

upvoted 2 times

RicardoD 1 year, 3 months ago

Selected Answer: A

A is correct

You can change the log data retention setting for CloudWatch logs. By default, logs are kept indefinitely and never expire. You can adjust the retention policy for each log group, keeping the indefinite retention, or choosing a retention period between 10 years and one day.

<https://docs.aws.amazon.com/managedservices/latest/userguide/log-customize-retention.html>

upvoted 1 times

jerryI 2 months, 2 weeks ago

you need random access after 90 days but option A is removing the logs

upvoted 1 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: D

Use Amazon S3 as the target. Enable an S3 Lifecycle policy to transition the logs to S3 Standard-Infrequent Access (S3 Standard-IA) after 90 days

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

D. Use Amazon S3 as the target. Enable an S3 Lifecycle policy to transition the logs to S3 Standard-Infrequent Access (S3 Standard-IA) after 90 days.

upvoted 2 times

animefan1 1 year, 11 months ago

Selected Answer: D

S3 will store logs. With life cycle, we can move it to different class. With Option A, log groups expiration will simply remove the logs and failing the 2nd request in question

upvoted 3 times

markw92 1 year, 11 months ago

A doesn't solve "90 days and then accessed intermittently" this statement. It sets expire after 90. Not sure otherwise A seems to be right choice since you can create dashboards etc.

upvoted 2 times

Bmarodi 2 years ago

Selected Answer: A

Option A meets these requirements.

upvoted 1 times

pentium75 1 year, 5 months ago

"Expiration of 90 days", but you need to access the log AFTER 90 days, just "intermittently".

upvoted 4 times

ocbn3wby 2 years, 4 months ago

Selected Answer: D

There's a table here that specifies that VPC Flow logs can go directly to S3. Does not need to go via CloudTrail and then to S3.

Nor via CW.

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/AWS-logs-and-resource-policy.html#AWS-logs-infrastructure-S3>

upvoted 4 times

techhb 2 years, 4 months ago

Selected Answer: D

we need to preserve logs hence D

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/CloudWatchLogsConcepts.html>

upvoted 3 times

mp165 2 years, 4 months ago

Selected Answer: D

D...agree that retention is the key word

upvoted 3 times

swolfgang 2 years, 4 months ago

Selected Answer: D

a is not, retention means delete after 90 days but questions say rarely access.

upvoted 3 times

mhmt4438 2 years, 4 months ago

Selected Answer: D

D. Use Amazon S3 as the target. Enable an S3 Lifecycle policy to transition the logs to S3 Standard-Infrequent Access (S3 Standard-IA) after 90 days.

By using Amazon S3 as the target for the VPC Flow Logs, the logs can be easily stored and accessed by the security team. Enabling an S3 Lifecycle policy to transition the logs to S3 Standard-Infrequent Access (S3 Standard-IA) after 90 days will automatically move the logs to a storage class that is optimized for infrequent access, reducing the storage costs for the company. The security team will still be able to access the logs as needed, even after they have been transitioned to S3 Standard-IA, but the storage cost will be optimized.

upvoted 5 times

laicos 2 years, 4 months ago

Selected Answer: D

I prefer D

"accessed intermittently" need logs after 90 days

upvoted 2 times

Parsons 2 years, 4 months ago

Selected Answer: D

No, D should be is correct.

"The logs will be frequently accessed for 90 days and then accessed intermittently." => We still need to store instead of deleting as the answer A.

upvoted 3 times

Aninina 2 years, 4 months ago

Selected Answer: D

D looks correct. This will meet the requirements of frequently accessing the logs for the first 90 days and then intermittently accessing them after that. S3 standard-IA is a storage class that is less expensive than S3 standard for infrequently accessed data, so it would be a more cost-effective option for storing the logs after the first 90 days.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #251

Topic 1


An Amazon EC2 instance is located in a private subnet in a new VPC. This subnet does not have outbound internet access, but the EC2 instance needs the ability to download monthly security updates from an outside vendor.

What should a solutions architect do to meet these requirements?

- A. Create an internet gateway, and attach it to the VPC. Configure the private subnet route table to use the internet gateway as the default route.
- B. Create a NAT gateway, and place it in a public subnet. Configure the private subnet route table to use the NAT gateway as the default route. **Most Voted**
- C. Create a NAT instance, and place it in the same subnet where the EC2 instance is located. Configure the private subnet route table to use the NAT instance as the default route.
- D. Create an internet gateway, and attach it to the VPC. Create a NAT instance, and place it in the same subnet where the EC2 instance is located. Configure the private subnet route table to use the internet gateway as the default route.

Correct Answer: B

Community vote distribution



Comments

mhmt4438 **Highly Voted** 1 year, 10 months ago

Selected Answer: B

B. Create a NAT gateway, and place it in a public subnet. Configure the private subnet route table to use the NAT gateway as the default route.

This approach will allow the EC2 instance to access the internet and download the monthly security updates while still being located in a private subnet. By creating a NAT gateway and placing it in a public subnet, it will allow the instances in the private subnet to access the internet through the NAT gateway. And then, configure the private subnet route table to use the NAT gateway as the default route. This will ensure that all outbound traffic is directed through the NAT gateway, allowing the EC2 instance to access the internet while still maintaining the security of the private subnet.

upvoted 10 times

Manjunathkb 1 year, 7 months ago

NAT gateway does not allow internet on it's own. It needs internet gateway too. None of the answers make sense

NAT gateway does not allow internet on its own. It needs internet gateway too. None of the answers make sense
upvoted 10 times

pentium75 11 months, 2 weeks ago

B says "place it in a public subnet", a public subnet needs an Internet Gateway so that is included in the answer.
upvoted 3 times

Manjunathkb 1 year, 7 months ago

refer below link
<https://aws.amazon.com/about-aws/whats-new/2021/06/aws-removes-nat-gateways-dependence-on-internet-gateway-for-private-communications/>
upvoted 3 times

TOR_0511 1 year ago

lol, thats for 'private connections'
upvoted 1 times

Uzbekistan Most Recent 8 months, 2 weeks ago

Selected Answer: B

B. Create a NAT gateway, and place it in a public subnet. Configure the private subnet route table to use the NAT gateway as the default route.

Explanation:

NAT Gateway: NAT (Network Address Translation) gateway is a managed service provided by AWS that allows EC2 instances in private subnets to access the internet while preventing inbound traffic from directly accessing them. You place the NAT gateway in a public subnet with an associated internet gateway, allowing it to send traffic to the internet.

Private Subnet Route Table: Configure the route table of the private subnet to route all outbound traffic (0.0.0.0/0) through the NAT gateway. This allows instances in the private subnet to access the internet through the NAT gateway while maintaining their private IP addresses and security.

upvoted 2 times

pentium75 11 months, 2 weeks ago

Selected Answer: B

A - if you "configure the private subnet route table to use the internet gateway" then it's no longer a private subnet
B - Correct (you place NAT GW in a public subnet and add it to the private subnet's route table)
C - NAT instance is deprecated, and it would still in a private subnet where it doesn't have Internet access
D - NAT instance is deprecated, and in that answer it is created but not even used
upvoted 3 times

EtherealBagel 1 year ago

yes, the nat gateway on its own does not allow connection to the internet. But the question specifies that it has been placed in a public subnet. public subnets are public because they have access to the internet via an internet gateway.
upvoted 2 times

xdkonorek2 1 year ago

Selected Answer: C

<https://docs.aws.amazon.com/vpc/latest/userguide/configure-subnets.html>

Public subnet – The subnet has a direct route to an internet gateway. Resources in a public subnet can access the public internet.

Private subnet – The subnet does not have a direct route to an internet gateway. Resources in a private subnet require a NAT device to access the public internet.

Both B and C have caveats but are both viable:

C - NAT Instance is used as a NAT device instead of NAT gateway, but it's still viable option

B - Have 2 redundant components - IGW and public subnet, and NAT gateway still would route traffic to IGW, and if VPC is a custom VPC routing has to be set up

upvoted 1 times

pentium75 11 months, 2 weeks ago

"NAT instance in the same subnet where the EC2 instance is located". how would you "use the NAT instance as the default

...first instance in the same subnet where the EC2 instance is located, then would you use the first instance as the default route" when it's in the same subnet?

upvoted 2 times

oluolope 1 year, 1 month ago

Selected Answer: D

A NAT Gateway should have one interface in each network it is connected to. I don't understand what it means when they say it is located either in the private or in the public network. It should be in both. Therefore, B and D do not really make sense. I choose D over B because there is a requirement to access the internet and although it is possible for the NAT to exist without an internet gateway, the later is still needed when internet access is required which is the case in this scenario.

upvoted 1 times

pentium75 11 months, 2 weeks ago

NAT Gateway must be in a public subnet as it needs Internet access. It can be specified in a private subnet's route table as a destination.

D doesn't make sense because you created an (outdated) NAT instance but don't use it (you point the route table to the Internet Gateway).

upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: B

Internet Gateway is required anyway to access the internet.

Option B makes more sense: Create a NAT gateway, and place it in a public subnet. Configure the private subnet route table to use the NAT gateway as the default route.

upvoted 2 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: B

B. Create a NAT gateway, and place it in a public subnet. Configure the private subnet route table to use the NAT gateway as the default route.

upvoted 2 times

cookieMr 1 year, 5 months ago

A. provides direct internet access to the private subnet, which is not desired in this case as the goal is to restrict outbound internet access.

B. allows the EC2 in the private subnet to access the internet through the NAT gateway, which acts as a proxy. It provides controlled outbound internet access while maintaining the security of the private subnet.

C. is similar to using a NAT gateway, but it involves using a NAT instance. NAT instances require more manual configuration and management compared to NAT gateways, making them a less preferred option.

D. combines the use of an internet gateway and a NAT instance, which is not necessary. It introduces unnecessary complexity and adds a NAT instance that requires additional management.

Overall, option B is the most appropriate solution as it utilizes a NAT gateway placed in a public subnet to enable controlled outbound internet access for the EC2 instance in the private subnet.

NAT Gateways are preferred over NAT Instances by AWS and in general.

upvoted 4 times

Bmarodi 1 year, 6 months ago

Selected Answer: B

Option B meets the requirements, hence B is right choice.

upvoted 2 times

Manjunathkb 1 year, 7 months ago

D would have been the answer if NAT gateway is installed in public subnet and not where EC2 is located. None of the answers are correct.

upvoted 2 times

AlessandraSAA 1 year, 9 months ago

why not C?

upvoted 1 times

UnluckyDucky 1 year, 8 months ago

Because NAT Gateways are preferred over NAT Instances by AWS and in general.

I have yet to find a situation where a NAT Instance would be more applicable than NAT Gateway which is fully managed and is overall an easier solution to implement - both in AWS questions or the real world.

upvoted 3 times

TungPham 1 year, 10 months ago

Selected Answer: B

Require NAT gateway

upvoted 2 times

techhb 1 year, 10 months ago

Selected Answer: B

Answer explained here <https://medium.com/@tshemku/aws-internet-gateway-vs-nat-gateway-vs-nat-instance-30523096df22>

upvoted 2 times

techhb 1 year, 10 months ago

Selected Answer: B

NAT Gateway is right choice

upvoted 2 times

bamishr 1 year, 11 months ago

Selected Answer: B

<https://www.examttopics.com/discussions/amazon/view/59966-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #252

Topic 1

A solutions architect needs to design a system to store client case files. The files are core company assets and are important. The number of files will grow over time.

The files must be simultaneously accessible from multiple application servers that run on Amazon EC2 instances. The solution must have built-in redundancy.

Which solution meets these requirements?

- A. Amazon Elastic File System (Amazon EFS) **Most Voted**
- B. Amazon Elastic Block Store (Amazon EBS)
- C. Amazon S3 Glacier Deep Archive
- D. AWS Backup

Correct Answer: A

Community vote distribution

A (100%)

Comments

Aninina **Highly Voted** 1 year, 4 months ago

Selected Answer: A

EFS Amazon Elastic File System (EFS) automatically grows and shrinks as you add and remove files with no need for management or provisioning.

upvoted 8 times

TungPham **Highly Voted** 1 year, 4 months ago

Selected Answer: A

"accessible from multiple application servers that run on Amazon EC2 instances"

upvoted 6 times

Chiquitabandita **Most Recent** 8 months, 1 week ago

Selected Answer: A

my choice is A but I think a better alternative would be S3 standard if offered wouldn't it be?

my choice is A but I think a better alternative would be S3 standard if offered wouldn't it be?
upvoted 5 times

TariqKipkemei 8 months, 2 weeks ago

Selected Answer: A

File system, scalable, multiple access = Amazon Elastic File System (Amazon EFS)
upvoted 3 times

Guru4Cloud 9 months ago

Selected Answer: A

Amazon Elastic File System (Amazon EFS)
upvoted 2 times

cookieMr 11 months, 2 weeks ago

Selected Answer: A

EFS provides a scalable and fully managed file storage service that can be accessed concurrently from multiple EC2. It offers built-in redundancy by storing data across multiple AZs within a region. With EFS, the client case files can be accessed by multiple application servers simultaneously, ensuring high availability and scalability as the number of files grows over time.

Option B, EBS, is a block-level storage service that is typically used for attaching to individual EC2 and does not provide concurrent access to multiple instances, making it unsuitable for this scenario.

Option C, S3 Glacier Deep Archive, is a long-term archival storage service and may not be suitable for active file access and simultaneous access from multiple application servers.

Option D, AWS Backup, is a centralized backup management service and does not provide the required simultaneous file access and redundancy features.

Therefore, the most suitable solution is Amazon EFS (option A).
upvoted 5 times

Bmarodi 1 year ago

Selected Answer: A

Option A meets the requirements, hence A is correct answer.
upvoted 2 times

moiraqi 1 year ago

What does "The solution must have built-in redundancy" mean
upvoted 1 times

KZM 1 year, 3 months ago

If the application servers are running on Linux or UNIX operating systems, EFS is a the most suitable solution for the given requirements.
upvoted 2 times

mhmt4438 1 year, 4 months ago

Selected Answer: A

Correct answer is A
upvoted 3 times

bamishr 1 year, 4 months ago

Selected Answer: A

<https://www.examttopics.com/discussions/amazon/view/68833-exam-aws-certified-solutions-architect-associate-saa-c02/>
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #253

Topic 1

A solutions architect has created two IAM policies: Policy1 and Policy2. Both policies are attached to an IAM group.

Policy 1

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "iam:Get*",
        "iam:List*",
        "kms:List*",
        "ec2:*",
        "ds:*",
        "logs:Get*",
        "logs:Describe*"
      ],
      "Resource": "*"
    }
  ]
}
```

Policy 2

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Deny",
      "Action": "ds:Delete*",
      "Resource": "*"
    }
  ]
}
```

A cloud engineer is added as an IAM user to the IAM group. Which action will the cloud engineer be able to perform?


- A. Deleting IAM users
- B. Deleting directories

C. Deleting Amazon EC2 instances **Most Voted**

D. Deleting logs from Amazon CloudWatch Logs

Correct Answer: C

Community vote distribution



Comments

JayBee65 **Highly Voted** 2 years, 4 months ago

ec2:* Allows full control of EC2 instances, so C is correct

The policy only grants get and list permission on IAM users, so not A
ds:Delete deny denies delete-directory, so not B, see
<https://awscli.amazonaws.com/v2/documentation/api/latest/reference/ds/index.html>
The policy only grants get and describe permission on logs, so not D
upvoted 23 times

vincent2023 4 months, 3 weeks ago

ds is related to directory service, don't understand how is it related to deleting ec2 instances ??
upvoted 1 times

mwwt2022 1 year, 5 months ago

great explanation
upvoted 2 times

Morinator **Highly Voted** 2 years, 4 months ago

Selected Answer: C

Explicite deny on directories, only available action for deleting is EC2
upvoted 7 times

surajkrishnamurthy **Most Recent** 4 months, 2 weeks ago

Selected Answer: C

In Policy number 2 its explicit deny to deleting directories so the IAM user wont be able to delete any directories
However in policy number 1 the statement "ec2:*", means the IAM user can perform all activities under the ec2
Therefore as per the given answer choice IAM user can delete an ec2 instance
so the choice is C: Deleting Amazon EC2 instances
upvoted 1 times

1e22522 10 months, 1 week ago

Selected Answer: B

This question and answers feel like a fever dream. what the hell is happening.
upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: C

Deleting Amazon EC2 instances
upvoted 2 times

Aninina 2 years, 4 months ago

Selected Answer: C

C : Deleting Amazon EC2 instances
upvoted 2 times

mhmt4438 2 years, 4 months ago

Selected Answer: C

Answer is C
upvoted 3 times

Aninina 2 years, 4 months ago

C : Deleting Amazon EC2 instances
upvoted 2 times

bamishr 2 years, 4 months ago

Selected Answer: C

<https://www.examtopics.com/discussions/amazon/view/27873-exam-aws-certified-solutions-architect-associate-saa-c02/>
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #254

Topic 1

A company is reviewing a recent migration of a three-tier application to a VPC. The security team discovers that the principle of least privilege is not being applied to Amazon EC2 security group ingress and egress rules between the application tiers.

What should a solutions architect do to correct this issue?

- A. Create security group rules using the instance ID as the source or destination.
- B. Create security group rules using the security group ID as the source or destination. **Most Voted**
- C. Create security group rules using the VPC CIDR blocks as the source or destination.
- D. Create security group rules using the subnet CIDR blocks as the source or destination.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Aninina **Highly Voted** 1 year, 10 months ago

Selected Answer: B

B. Create security group rules using the security group ID as the source or destination.

This way, the security team can ensure that the least privileged access is given to the application tiers by allowing only the necessary communication between the security groups. For example, the web tier security group should only allow incoming traffic from the load balancer security group and outgoing traffic to the application tier security group. This approach provides a more granular and secure way to control traffic between the different tiers of the application and also allows for easy modification of access if needed.

It's also worth noting that it's good practice to minimize the number of open ports and protocols, and use security groups as a first line of defense, in addition to network access control lists (ACLs) to control traffic between subnets.

upvoted 13 times

Wael216 **Highly Voted** 1 year, 9 months ago

Selected Answer: B

By using security group IDs, the ingress and egress rules can be restricted to only allow traffic from the necessary source or destination, and to deny all other traffic. This ensures that only the minimum required traffic is allowed between the application tiers.

Option A is not the best choice because using the instance ID as the source or destination would allow traffic from any instance with that ID, which may not be limited to the specific application tier.

with that ID, which may not be limited to the specific application tier.

Option C is also not the best choice because using VPC CIDR blocks would allow traffic from any IP address within the VPC, which may not be limited to the specific application tier.

Option D is not the best choice because using subnet CIDR blocks would allow traffic from any IP address within the subnet, which may not be limited to the specific application tier.

upvoted 9 times

Guru4Cloud Most Recent 1 year, 2 months ago

Selected Answer: B

Create security group rules using the security group ID as the source or destination.

This way, the security team can ensure that the least privileged access is given to the application tiers by allowing only the necessary communication between the security groups. For example, the web tier security group should only allow incoming traffic from the load balancer security group and outgoing traffic to the application tier security group. This approach provides a more granular and secure way to control traffic between the different tiers of the application and also allows for easy modification of access if needed.

It's also worth noting that it's good practice to minimize the number of open ports and protocols, and use security groups as a first line of defense, in addition to network access control lists (ACLs) to control traffic between subnets.

upvoted 3 times

cookieMr 1 year, 5 months ago

Selected Answer: B

A. would limit the traffic based on specific instances, which may not be the most suitable solution for applying the principle of least privilege between application tiers.

B. By using security group IDs in the rules, you can precisely control the traffic between application tiers, allowing only the necessary communication and adhering to the principle of least privilege.

C. would apply broad rules based on the entire VPC CIDR blocks, which may not provide the necessary level of granularity required for secure communication between specific application tiers.

D. would limit the traffic based on subnet CIDR blocks, which may not be sufficient for ensuring proper security between application tiers.

In summary, using security group IDs (Option B) is the recommended approach as it allows for precise control of traffic between application tiers, aligning with the principle of least privilege.

upvoted 7 times

foha2012 10 months, 3 weeks ago

with option A. How would you use instance ID in security group inbound rules ?

upvoted 2 times

Bmarodi 1 year, 6 months ago

Selected Answer: B

I vote for option B.

upvoted 2 times

LuckyAro 1 year, 10 months ago

Selected Answer: B

. Create security group rules using the security group ID as the source or destination

upvoted 2 times

techhb 1 year, 10 months ago

Security Group Rules apply to instances

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/security-group-rules.html>

upvoted 2 times

mhmt4438 1 year, 10 months ago

Selected Answer: B

Correct answer is B

upvoted 3 times

bamishr 1 year, 11 months ago

Selected Answer: B

<https://www.examttopics.com/discussions/amazon/view/46463-exam-aws-certified-solutions-architect-associate-saa-c02/>
upvoted 2 times

Morinator 1 year, 11 months ago

Selected Answer: B

B right
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/security-group-rules.html>
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #255

Topic 1

A company has an ecommerce checkout workflow that writes an order to a database and calls a service to process the payment. Users are experiencing timeouts during the checkout process. When users resubmit the checkout form, multiple unique orders are created for the same desired transaction.

How should a solutions architect refactor this workflow to prevent the creation of multiple orders?

- A. Configure the web application to send an order message to Amazon Kinesis Data Firehose. Set the payment service to retrieve the message from Kinesis Data Firehose and process the order.
- B. Create a rule in AWS CloudTrail to invoke an AWS Lambda function based on the logged application path request. Use Lambda to query the database, call the payment service, and pass in the order information.
- C. Store the order in the database. Send a message that includes the order number to Amazon Simple Notification Service (Amazon SNS). Set the payment service to poll Amazon SNS, retrieve the message, and process the order.
- D. Store the order in the database. Send a message that includes the order number to an Amazon Simple Queue Service (Amazon SQS) FIFO queue. Set the payment service to retrieve the message and process the order. Delete the message from the queue. **Most Voted**

Correct Answer: D*Community vote distribution*

D (100%)

Comments**Aninina** **Highly Voted** 1 year, 10 months ago**Selected Answer: D**

D. Store the order in the database. Send a message that includes the order number to an Amazon Simple Queue Service (Amazon SQS) FIFO queue. Set the payment service to retrieve the message and process the order. Delete the message from the queue.

This approach ensures that the order creation and payment processing steps are separate and atomic. By sending the order information to an SQS FIFO queue, the payment service can process the order one at a time and in the order they were received. If the payment service is unable to process an order, it can be retried later, preventing the creation of multiple orders. The deletion of the message from the queue after it is processed will prevent the same message from being processed multiple times.

It's worth noting that FIFO queues guarantee that messages are processed in the order they are received, and prevent duplicates.

upvoted 11 times

cookieMr Highly Voted 1 year, 5 months ago

Selected Answer: D

A. is not a suitable solution for preventing the creation of multiple orders. This approach does not guarantee the sequential and reliable processing of orders.

B. is not an appropriate solution for preventing the creation of multiple orders. CloudTrail is primarily used for logging and auditing API activity, and invoking a Lambda based on the logged request does not ensure the correct order processing.

C. is not a suitable solution. SNS is a publish-subscribe messaging service, and polling it may result in delayed processing and potential order duplication.

D. is the correct solution. Using an SQS FIFO ensures that the orders are processed in a sequential and reliable manner, preventing the creation of multiple orders for the same transaction.

upvoted 7 times

TariqKipkemei Most Recent 1 year, 2 months ago

Selected Answer: D

if the backend can not keep up, queue the tasks.

upvoted 2 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: D

D. Store the order in the database. Send a message that includes the order number to an Amazon Simple Queue Service (Amazon SQS) FIFO queue. Set the payment service to retrieve the message and process the order. Delete the message from the queue.

upvoted 2 times

animefan1 1 year, 5 months ago

Selected Answer: D

The question is related in breaking down the flow. SQS is go-to choice to decouple & DB will be used to store

upvoted 2 times

antropaws 1 year, 5 months ago

Why not A?

upvoted 1 times

pentium75 11 months, 2 weeks ago

Because Kinesis Data Firehose is for ingestion of streaming data, not queuing items.

upvoted 4 times

Wael216 1 year, 9 months ago

Selected Answer: D

The use of a FIFO queue in Amazon SQS ensures that messages are processed in the order they are received.

upvoted 2 times

mhmt4438 1 year, 10 months ago

Selected Answer: D

<https://www.examttopics.com/discussions/amazon/view/95026-exam-aws-certified-solutions-architect-associate-saa-c03/>

upvoted 4 times

bamishr 1 year, 11 months ago

Selected Answer: D

asnwer is d

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #256

Topic 1

A solutions architect is implementing a document review application using an Amazon S3 bucket for storage. The solution must prevent accidental deletion of the documents and ensure that all versions of the documents are available. Users must be able to download, modify, and upload documents.

Which combination of actions should be taken to meet these requirements? (Choose two.)

- A. Enable a read-only bucket ACL.
- B. Enable versioning on the bucket. **Most Voted**
- C. Attach an IAM policy to the bucket.
- D. Enable MFA Delete on the bucket. **Most Voted**
- E. Encrypt the bucket using AWS KMS.

Correct Answer: BD

Community vote distribution

BD (100%)

Comments

LoXoL 10 months, 2 weeks ago

Selected Answer: BD

No Brainer: B & D

upvoted 3 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: BD

Prevent accidental deletion of the documents = Enable MFA Delete on the bucket

Ensure that all versions of the documents are available = Enable versioning on the bucket

upvoted 4 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: BD

Options B & D are the correct answers.

upvoted 2 times

cookieMr 1 year, 5 months ago

Selected Answer: BD

B. allows multiple versions of objects in the S3 bucket to be stored. This ensures that all versions of the documents are available, even if they are accidentally overwritten or deleted.

D. adds an extra layer of protection against accidental deletion of objects in the bucket. With MFA Delete enabled, a user would need to provide an additional authentication factor to successfully delete objects from the bucket. This helps prevent accidental or unauthorized deletions and provides an extra level of security for critical documents.

A. would restrict users from modifying or uploading documents. It would not meet the requirement of allowing users to download, modify, and upload documents.

C. can control access permissions to the bucket, it does not specifically address the requirement of preventing accidental deletion or ensuring availability of all versions of the documents.

E. Encryption focuses on data protection rather than versioning and deletion prevention.

upvoted 4 times

Bmarodi 1 year, 6 months ago

Selected Answer: BD

Options B & D are the correct answers.

upvoted 2 times

Wael216 1 year, 9 months ago

Selected Answer: BD

no doubts

upvoted 3 times

MinHyeok 1 year, 10 months ago

□□□ □□□□□□

upvoted 4 times

akdavsan 1 year, 10 months ago

Selected Answer: BD

b and d ofc

upvoted 2 times

LuckyAro 1 year, 10 months ago

Selected Answer: BD

B & D Definitely.

upvoted 2 times

david76x 1 year, 10 months ago

Selected Answer: BD

B & D is correct

upvoted 2 times

Aninina 1 year, 10 months ago

Selected Answer: BD

B and D for sure guys

upvoted 3 times

mhmt4438 1 year, 10 months ago

Selected Answer: BD

<https://www.examttopics.com/discussions/amazon/view/21969-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #257

Topic 1


A company is building a solution that will report Amazon EC2 Auto Scaling events across all the applications in an AWS account. The company needs to use a serverless solution to store the EC2 Auto Scaling status data in Amazon S3. The company then will use the data in Amazon S3 to provide near-real-time updates in a dashboard. The solution must not affect the speed of EC2 instance launches.

How should the company move the data to Amazon S3 to meet these requirements?

- A. Use an Amazon CloudWatch metric stream to send the EC2 Auto Scaling status data to Amazon Kinesis Data Firehose. Store the data in Amazon S3. **Most Voted**
- B. Launch an Amazon EMR cluster to collect the EC2 Auto Scaling status data and send the data to Amazon Kinesis Data Firehose. Store the data in Amazon S3.
- C. Create an Amazon EventBridge rule to invoke an AWS Lambda function on a schedule. Configure the Lambda function to send the EC2 Auto Scaling status data directly to Amazon S3.
- D. Use a bootstrap script during the launch of an EC2 instance to install Amazon Kinesis Agent. Configure Kinesis Agent to collect the EC2 Auto Scaling status data and send the data to Amazon Kinesis Data Firehose. Store the data in Amazon S3.

Correct Answer: A

Community vote distribution



Comments

pentium75 **Highly Voted** 1 year, 5 months ago

Selected Answer: A

B - EMR cluster is for Big Data, has nothing to do with this

C - invokes the function "on a schedule", but you want to capture events

D - Could work, but would be overcomplex and would "affect the speed of EC2 instance launches" (which it should not)

upvoted 10 times

LoXoL 1 year, 4 months ago

Right.

upvoted 2 times

Guru4Cloud Highly Voted 1 year, 9 months ago

Selected Answer: A

This solution meets the requirements because it is serverless and does not affect the speed of EC2 instance launches. Amazon CloudWatch metric streams can continuously stream CloudWatch metrics to destinations such as Amazon S3. Amazon Kinesis Data Firehose can capture, transform, and deliver streaming data into data lakes, data stores, and analytics services. It can directly put the data into Amazon S3, which can then be used for near-real-time updates in a dashboard.

upvoted 8 times

reviewmine Most Recent 1 year, 3 months ago

Selected Answer: A

Answer A: Near real time --> Amazon Kinesis Data Firehose

upvoted 5 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: A

You can use metric streams to continually stream CloudWatch metrics to a destination of your choice, with near-real-time delivery and low latency. Supported destinations include AWS destinations such as Amazon Simple Storage Service and several third-party service provider destinations.

Main usage scenarios for CloudWatch metric streams: Data lake— Create a metric stream and direct it to an Amazon Kinesis Data Firehose delivery stream that delivers your CloudWatch metrics to a data lake such as Amazon S3.

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CloudWatch-Metric-Streams.html#:~:text=CloudWatch%20metric%20streams>

upvoted 4 times

Valder21 1 year, 9 months ago

Selected Answer: C

Kinesis is for data streams not events. So, C

upvoted 1 times

pentium75 1 year, 5 months ago

C invokes the Lambda function "on a schedule". It would collect the scaling status during its runs. But you don't want the hourly status, you want to report "scaling events".

upvoted 2 times

cookieMr 1 year, 11 months ago

Selected Answer: A

B. introduces unnecessary complexity and overhead for collecting and sending the EC2 Auto Scaling status data to S3. It is not the most efficient serverless solution for this specific requirement.

C. would introduce delays in data updates, as it is not triggered in real-time. Additionally, it adds unnecessary overhead and complexity compared to using a direct data stream.

D. introduces additional dependencies and management overhead. It may also impact the speed of EC2 instance launches, which is a requirement that needs to be avoided.

Overall, option A provides a streamlined and serverless solution by leveraging CloudWatch metric streams and Kinesis Data Firehose to efficiently capture and store the EC2 Auto Scaling status data in S3 without affecting the speed of EC2 instance launches.

upvoted 6 times

markw92 1 year, 11 months ago

A: I was thinking D is the answer but the solution should not impact ec2 launches will make the difference and i fast read the question. A is a right choice.

upvoted 2 times

Rahulbit34 2 years, 1 month ago

A because of near real time scenario

upvoted 4 times

UnluckyDucky 2 years, 2 months ago

Selected Answer: C

Both A and C are applicable - no doubt there.

C is more straightforward and to the point of the question imho.

upvoted 3 times

UnluckyDucky 2 years, 2 months ago

Changing my answer to *A* as the dashboard will provide near-real updates.

Unless the lambda is configured to run every minute which is not common with schedules - it is not considered near real-time.

upvoted 4 times

JohnYu 8 months ago

While EventBridge can capture events, scheduling Lambda functions to poll data is less efficient and may introduce latency, which would not meet the near-real-time requirement.

upvoted 2 times

bdp123 2 years, 3 months ago

Selected Answer: A

Serverless solution and near real time

upvoted 4 times

Stanislav4907 2 years, 3 months ago

Selected Answer: A

near real time -eliminates c

upvoted 2 times

aakashkumar1999 2 years, 4 months ago

Selected Answer: A

Answer is A

upvoted 2 times

devonwho 2 years, 4 months ago

Selected Answer: A

You can use metric streams to continually stream CloudWatch metrics to a destination of your choice, with near-real-time delivery and low latency. One of the use cases is Data Lake: create a metric stream and direct it to an Amazon Kinesis Data Firehose delivery stream that delivers your CloudWatch metrics to a data lake such as Amazon S3.

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CloudWatch-Metric-Streams.html>

upvoted 3 times

Stanislav4907 2 years, 4 months ago

Selected Answer: A

Option C, using an Amazon EventBridge rule to invoke an AWS Lambda function on a schedule to send the EC2 Auto Scaling status data directly to Amazon S3, may not be the best choice because it may not provide real-time updates to the dashboard.

A schedule-based approach with an EventBridge rule and Lambda function may not be able to deliver the data in near real-time, as the EC2 Auto Scaling status data is generated dynamically and may not always align with the schedule set by the EventBridge rule.

Additionally, using a schedule-based approach with EventBridge and Lambda also has the potential to create latency, as there may be a delay between the time the data is generated and the time it is sent to S3.

In this scenario, using Amazon CloudWatch and Kinesis Data Firehose as described in Option A, provides a more reliable and near real-time solution.

upvoted 2 times

MikelH93 2 years, 4 months ago

Selected Answer: A

A seems to be the right answer. Don't think C could be correct as it says "near real-time" and C is on schedule
upvoted 2 times

KAUS2 2 years, 4 months ago

Selected Answer: C

C. Create an Amazon EventBridge rule to invoke an AWS Lambda function on a schedule. Configure the Lambda function to send the EC2 Auto Scaling status data directly to Amazon S3.

upvoted 2 times

pentium75 1 year, 5 months ago

"On a schedule" but you want to capture events, not a regular status report.

upvoted 2 times

techhb 2 years, 4 months ago

Selected Answer: A

A seems right choice but serverless keyword confuses, and cloud watch metric stream is serverless too.

upvoted 3 times

pentium75 1 year, 5 months ago

But A is serverless.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #258

Topic 1


A company has an application that places hundreds of .csv files into an Amazon S3 bucket every hour. The files are 1 GB in size. Each time a file is uploaded, the company needs to convert the file to Apache Parquet format and place the output file into an S3 bucket.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an AWS Lambda function to download the .csv files, convert the files to Parquet format, and place the output files in an S3 bucket. Invoke the Lambda function for each S3 PUT event.
- B. Create an Apache Spark job to read the .csv files, convert the files to Parquet format, and place the output files in an S3 bucket. Create an AWS Lambda function for each S3 PUT event to invoke the Spark job.
- C. Create an AWS Glue table and an AWS Glue crawler for the S3 bucket where the application places the .csv files. Schedule an AWS Lambda function to periodically use Amazon Athena to query the AWS Glue table, convert the query results into Parquet format, and place the output files into an S3 bucket.
- D. Create an AWS Glue extract, transform, and load (ETL) job to convert the .csv files to Parquet format and place the output files into an S3 bucket. Create an AWS Lambda function for each S3 PUT event to invoke the ETL job. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

Parsons **Highly Voted** 2 years, 4 months ago

Selected Answer: D

No, D should be correct.

"LEAST operational overhead" => Should you fully manage service like Glue instead of manually like the answer A.
upvoted 16 times

awsgeek75 1 year, 5 months ago

I also think it's D but remember that D requires writing ETL logic in AWS Glue (nothing in question says how complex it will be). AWS Lambda for CSV could be simple also (imagine NodeJS and millions of libraries support or Python's parsing) so both could be operationally on par to each other. Logically D makes more sense but in practice, AWS Glue rarely works with out of

...can be operationally on par to each other, eg. using a makes more sense than in practice, and also rarely, none that out of the box ETL and becomes a maintenance overhead in itself.

upvoted 1 times

Guru4Cloud Highly Voted 1 year, 9 months ago

Selected Answer: D

This solution meets the requirements with the least operational overhead because AWS Glue is a fully managed ETL service that makes it easy to move data between data stores. AWS Glue can read .csv files from an S3 bucket and write the data into Parquet format in another S3 bucket. The AWS Lambda function can be triggered by an S3 PUT event when a new .csv file is uploaded, and it can start the AWS Glue ETL job to convert the file to Parquet format. This solution does not require managing any servers or clusters, which reduces operational overhead.

upvoted 5 times

TariqKipkemei Most Recent 1 year, 8 months ago

Selected Answer: D

AWS Glue can run your extract, transform, and load (ETL) jobs as new data arrives. For example, you can configure AWS Glue to initiate your ETL jobs to run as soon as new data becomes available in Amazon Simple Storage Service (S3). Clearly you don't need a lambda function to initiate the ETL job.

<https://aws.amazon.com/glue/#:~:text=to%20initiate%20your-,ETL,-jobs%20to%20run>

Option A requires writing code to perform the file conversion.

In the exam option D would be the best answer.

upvoted 4 times

cookieMr 1 year, 11 months ago

D is correct

upvoted 2 times

cookieMr 1 year, 11 months ago

A. introduces significant operational overhead. This approach requires managing the Lambda, handling concurrency, and ensuring proper error handling for large file sizes, which can be challenging.

B. adds unnecessary complexity and operational overhead. Managing the Spark job, handling scalability, and coordinating the Lambda invocations for each file upload can be cumbersome.

C. introduces additional complexity and may not be the most efficient solution. It involves managing Glue resources, scheduling Lambda, and querying data even when no new files are uploaded.

Option D leverages AWS Glue's ETL capabilities, allowing you to define and execute a data transformation job at scale. By invoking the ETL job using a Lambda function for each S3 PUT event, you can ensure that files are efficiently converted to Parquet format without the need for manual intervention. This approach minimizes operational overhead and provides a streamlined and scalable solution.

upvoted 5 times

F629 1 year, 11 months ago

Selected Answer: A

Both A and D can work, but A is more simple. It's more close to the "Least Operational effort".

upvoted 1 times

pentium75 1 year, 5 months ago

Creating, maintaining and supporting custom code that does the same as a ready-made serverless service is NEVER "least operational effort".

upvoted 2 times

pentium75 1 year, 5 months ago

Oh, and A can't handle 1 GB files.

upvoted 1 times

jaswantn 1 year, 4 months ago

Now Lambda support 1 GB to 10 GB.

upvoted 2 times

JA2018 6 months, 3 weeks ago

But what if the file conversion process exceeds more than 15 minutes for each file? How does Lambda fits into this picture?

upvoted 2 times

shanwford 2 years, 2 months ago

Selected Answer: D

The maximum size for a Lambda event payload is 256 KB - so (A) didn't work with 1GB Files. Glue is recommended for the Parquet Transformation of AWS.

upvoted 3 times

jennyka76 2 years, 3 months ago

ANS - d

<https://aws.amazon.com/blogs/database/how-to-extract-transform-and-load-data-for-analytic-processing-using-aws-glue-part-2/>

- READ ARTICLE -

upvoted 3 times

aws4myself 2 years, 4 months ago

Here A is the correct answer. The reason here is the least operational overhead.

A ==> S3 - Lambda - S3

D ==> S3 - Lambda - Glue - S3

Also, glue cannot convert on fly automatically, you need to write some code there. If you write the same code in lambda it will convert the same and push the file to S3

Lambda has max memory of 128 MB to 10 GB. So, it can handle it easily.

And we need to consider cost also, glue cost is more. Hope many from this forum realize these differences.

upvoted 5 times

JA2018 6 months, 3 weeks ago

But what if the file conversion process exceeds more than 15 minutes for each file? How does Lambda fits into this picture?

upvoted 1 times

JA2018 6 months, 3 weeks ago

IIRC, many folks on this forum mentioned that, rule of thumbs for AWS exams, opt for AWS managed services...

upvoted 1 times

LuckyAro 2 years, 4 months ago

We also need to stay with the question, cost was not a consideration in the question.

upvoted 2 times

nder 2 years, 3 months ago

Cost is not a factor. AWS Glue is a fully managed service therefore, it's the least operational overhead

upvoted 5 times

JayBee65 2 years, 4 months ago

A is unlikely to work as Lambda may struggle with 1GB size: "< 64 MB, beyond which lambda is likely to hit memory caps", see <https://stackoverflow.com/questions/41504095/creating-a-parquet-file-on-aws-lambda-function>

upvoted 3 times

jainparag1 2 years, 4 months ago

Should be D as Glue is self managed service and provides tel job for converting cab files to parquet off the shelf.

upvoted 2 times

Jxtat 2 years, 4 months ago

Selected Answer: D

<https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/three-aws-glue-etl-job-types-for-converting-data-to-apache-parquet.html>

upvoted 2 times

techhb 2 years, 4 months ago

AWS Glue is right solution here.

upvoted 2 times

mp165 2 years, 4 months ago

Selected Answer: D

I am thinking D.

A says lambda will download the .csv...but to where? that seem manual based on that

upvoted 2 times

mhmt4438 2 years, 4 months ago

Selected Answer: A

I think A

upvoted 1 times

bamishr 2 years, 4 months ago

Selected Answer: A

<https://www.examttopics.com/discussions/amazon/view/83201-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #259

Topic 1

A company is implementing new data retention policies for all databases that run on Amazon RDS DB instances. The company must retain daily backups for a minimum period of 2 years. The backups must be consistent and restorable.

Which solution should a solutions architect recommend to meet these requirements?

- A. Create a backup vault in AWS Backup to retain RDS backups. Create a new backup plan with a daily schedule and an expiration period of 2 years after creation. Assign the RDS DB instances to the backup plan. **Most Voted**
- B. Configure a backup window for the RDS DB instances for daily snapshots. Assign a snapshot retention policy of 2 years to each RDS DB instance. Use Amazon Data Lifecycle Manager (Amazon DLM) to schedule snapshot deletions.
- C. Configure database transaction logs to be automatically backed up to Amazon CloudWatch Logs with an expiration period of 2 years.
- D. Configure an AWS Database Migration Service (AWS DMS) replication task. Deploy a replication instance, and configure a change data capture (CDC) task to stream database changes to Amazon S3 as the target. Configure S3 Lifecycle policies to delete the snapshots after 2 years.

Correct Answer: A*Community vote distribution*

A (95%)

B (5%)

Comments**cookieMr** **Highly Voted** 1 year, 11 months ago**Selected Answer: A**

A. suggests using AWS Backup, a centralized backup management service, to retain RDS backups. A backup vault is created, and a backup plan is defined with a daily schedule and a 2-year retention period for backups. RDS DB instances are assigned to this backup plan.

B. it does not address the requirement for consistent and restorable backups. Snapshots are point-in-time backups and may not provide the desired level of consistency.

C. it is not designed to provide the backup and restore functionality required for databases. It does not ensure the backups are consistent or provide an easy restore mechanism.

D. it does not address the requirement for daily backups and retention of consistent backups. It focuses more on replication and change data capture rather than backup and restore.

and change data capture rather than backup and restore.

upvoted 11 times

bamishr Highly Voted 2 years, 4 months ago

Selected Answer: A

Create a backup vault in AWS Backup to retain RDS backups. Create a new backup plan with a daily schedule and an expiration period of 2 years after creation. Assign the RDS DB instances to the backup plan.

upvoted 5 times

Ojonugwa Most Recent 1 year, 2 months ago

The solution architect should recommend option B - Configure a backup window for the RDS DB instances for daily snapshots. Assign a snapshot retention policy of 2 years to each RDS DB instance. Use Amazon Data Lifecycle Manager (Amazon DLM) to schedule snapshot deletions.

This meets the requirements of:

Retaining daily backups for a minimum of 2 years. The daily snapshots captured within the backup window provide consistent, restorable backups on a daily basis.

Assigning a snapshot retention policy of 2 years ensures the snapshots are retained for the required period.

Using Amazon DLM allows automatically deleting snapshots older than 2 years to comply with the retention period in a cost-effective manner without manual administration.

The other option of using AWS Backup vault is not as suitable since it has limitations such as 35 day maximum retention for automated backups. Option B provides a native RDS solution capable of meeting the long term 2 year retention requirements.

upvoted 3 times

JA2018 6 months, 3 weeks ago

I thought, for AWS exams, as a rule of thumbs, we should opt for AWS managed services?

upvoted 1 times

kennylin689 9 months, 3 weeks ago

If not specified, AWS Backup Vault Lock will not enforce a maximum retention period. If specified, backup and copy jobs to this vault with lifecycle retention periods longer than the maximum retention period will fail. Recovery points already saved in the vault prior to the vault lock's creation are not affected. The longest maximum retention period you can specify is 36500 days (approximately 100 years).

upvoted 2 times

vijaykamal 1 year, 8 months ago

Selected Answer: B

Here's why Option B is the best choice:

Backup Window: Configuring a backup window for daily snapshots ensures that consistent backups are taken at the specified time each day. This helps maintain data integrity and consistency.

Snapshot Retention Policy: Assigning a snapshot retention policy of 2 years to each RDS DB instance ensures that the backups are retained for the required duration.

Amazon Data Lifecycle Manager (Amazon DLM): Amazon DLM can be used to automate the management of EBS snapshots, including RDS snapshots. You can configure Amazon DLM to schedule snapshot deletions, making it easier to manage the retention policy without manual intervention.

Option A (AWS Backup) is primarily used for managing backups of resources that may not have built-in backup capabilities, but for Amazon RDS, it's better to use the built-in snapshot capabilities and Amazon DLM for snapshot retention.

upvoted 2 times

pentium75 1 year, 5 months ago

"AWS Backup is primarily used for managing backups of resources that may not have built-in backup capabilities" says who?

upvoted 3 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: A

Create a backup vault in AWS Backup to retain RDS backups. Create a new backup plan with a daily schedule and an expiration

Create a backup vault in AWS Backup to retain RDS backups. Create a new backup plan with a daily schedule and an expiration period of 2 years after creation. Assign the RDS DB instances to the backup plan.

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

A. Create a backup vault in AWS Backup to retain RDS backups. Create a new backup plan with a daily schedule and an expiration period of 2 years after creation. Assign the RDS DB instances to the backup plan

upvoted 3 times

animefan1 1 year, 11 months ago

Selected Answer: A

Backups work with EBS, FSX, RDS. Its managed & can has vault option for more better control over backup retention

upvoted 4 times

markw92 1 year, 11 months ago

Why not B?

upvoted 2 times

_deepsi_dee29 2 years ago

Selected Answer: A

A is correct

upvoted 1 times

antropaws 2 years ago

Why not D?

Creating tasks for ongoing replication using AWS DMS: You can create an AWS DMS task that captures ongoing changes from the source data store. You can do this capture while you are migrating your data. You can also create a task that captures ongoing changes after you complete your initial (full-load) migration to a supported target data store. This process is called ongoing replication or change data capture (CDC). AWS DMS uses this process when replicating ongoing changes from a source data store.

upvoted 1 times

pentium75 1 year, 5 months ago

Requirement is backup, not migration.

upvoted 3 times

gold4otas 2 years, 2 months ago

Selected Answer: A

A. Create a backup vault in AWS Backup to retain RDS backups. Create a new backup plan with a daily schedule and an expiration period of 2 years after creation. Assign the RDS DB instances to the backup plan.

upvoted 2 times

techhb 2 years, 4 months ago

Selected Answer: A

A is right choice

upvoted 4 times

Aninina 2 years, 4 months ago

Selected Answer: A

A A A A A

upvoted 3 times

mhmt4438 2 years, 4 months ago

Selected Answer: A

Correct answer is A

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #260

Topic 1

A company's compliance team needs to move its file shares to AWS. The shares run on a Windows Server SMB file share. A self-managed on-premises Active Directory controls access to the files and folders.


The company wants to use Amazon FSx for Windows File Server as part of the solution. The company must ensure that the on-premises Active Directory groups restrict access to the FSx for Windows File Server SMB compliance shares, folders, and files after the move to AWS. The company has created an FSx for Windows File Server file system.

Which solution will meet these requirements?

- A. Create an Active Directory Connector to connect to the Active Directory. Map the Active Directory groups to IAM groups to restrict access.
- B. Assign a tag with a Restrict tag key and a Compliance tag value. Map the Active Directory groups to IAM groups to restrict access.
- C. Create an IAM service-linked role that is linked directly to FSx for Windows File Server to restrict access.
- D. Join the file system to the Active Directory to restrict access. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

mhmt4438 **Highly Voted** 2 years, 4 months ago

Selected Answer: D

D. Join the file system to the Active Directory to restrict access.

Joining the FSx for Windows File Server file system to the on-premises Active Directory will allow the company to use the existing Active Directory groups to restrict access to the file shares, folders, and files after the move to AWS. This option allows the company to continue using their existing access controls and management structure, making the transition to AWS more seamless.

upvoted 23 times

cookieMr **Highly Voted** 1 year, 11 months ago

Selected Answer: D

D. allows the file system to leverage the existing AD infrastructure for authentication and access control.

Option A is incorrect because mapping the AD groups to IAM groups is not applicable in this scenario. IAM is primarily used for managing access to AWS resources, while the requirement is to integrate with the on-premises AD for access control.

Option B is incorrect because assigning a tag with a Restrict tag key and a Compliance tag value does not provide the necessary integration with the on-premises AD for access control. Tags are used for organizing and categorizing resources and do not provide authentication or access control mechanisms.

Option C is incorrect because creating an IAM service-linked role linked directly to FSx for Windows File Server does not integrate with the on-premises AD. IAM roles are used within AWS for managing permissions and do not provide the necessary integration with external AD systems.

upvoted 9 times

zdi561 Most Recent 4 months, 3 weeks ago

Selected Answer: D

[https://aws.amazon.com/blogs/storage/using-amazon-fsx-for-windows-file-server-with-an-on-premises-active-directory/#:~:text=Perform%20the%20following%20steps%3A,shown%20in%20the%20following%20screenshot\).](https://aws.amazon.com/blogs/storage/using-amazon-fsx-for-windows-file-server-with-an-on-premises-active-directory/#:~:text=Perform%20the%20following%20steps%3A,shown%20in%20the%20following%20screenshot).)

upvoted 1 times

Lin878 1 year ago

Selected Answer: D

When you create a file system with Amazon FSx, you join it to your Active Directory domain to provide user authentication and file- and folder-level access control.

upvoted 4 times

bujuman 1 year, 3 months ago

Selected Answer: D

D is relevant and accurate answer when we consider this:

<https://docs.aws.amazon.com/fsx/latest/WindowsGuide/creating-joined-ad-file-systems.html>

"When you create a new FSx for Windows File Server file system, you can configure Microsoft Active Directory integration so that it joins to your self-managed Microsoft Active Directory domain. To do this, provide the following information for your Microsoft Active Directory"

upvoted 3 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

The on-premise AD already has restrictions via group in place so D makes no sense as the groups are already linked to file system.

"The company must ensure that the on-premises Active Directory groups restrict access to the FSx for Windows File Server SMB compliance shares, folders, and files after the move to AWS."

The question is about linking the on-prem permissions to the new FSx server on AWS and this can only be done by A

upvoted 1 times

awsgeek75 1 year, 4 months ago

Actually neither A nor D make sense.

"A self-managed on-premises Active Directory controls access to the files and folders." This makes D sound useless and at the same time does not allow the on-prem AD to control file access on FSx.

A uses IAM roles which is irrelevant to this setup.

BC are totally wrong

upvoted 1 times

meowruki 1 year, 6 months ago

Selected Answer: D

Option A: Creating an Active Directory Connector and mapping groups to IAM groups is more relevant for AWS Directory Service, such as AWS Managed Microsoft AD, and not for integrating with existing on-premises Active Directory.

Option B: Using tags is typically not used for access control purposes. Tags are metadata and are not directly involved in user

authentication and authorization.

Option C: Creating an IAM service-linked role directly linked to FSx for Windows File Server is not the standard approach for integrating with existing on-premises Active Directory.

upvoted 4 times

wrmari 1 year, 8 months ago

Selected Answer: D

<https://docs.aws.amazon.com/fsx/latest/WindowsGuide/self-managed-AD.html>

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

This allows the on-premises Active Directory to manage permissions to the FSx file shares, meeting the key requirement to use existing AD groups to control access after migrating to AWS.

Joining FSx to the AD domain allows the native file system permissions, users, and groups to be applied from Active Directory. Access is handled seamlessly via the trust relationship between FSx and AD.

The other options would not leverage the existing AD identities and groups

upvoted 4 times

Guru4Cloud 1 year, 9 months ago

The other options would not leverage the existing AD identities and groups:

A) AD Connector and IAM groups would require re-mapping AD groups to IAM, adding complexity. Native AD integration is simpler.

B) Tags and IAM groups also don't use native AD semantics.

C) Service-linked roles are not applicable for managing end user access.

So D is the correct option to meet the requirements using the native Active Directory integration built into FSx for Windows.

upvoted 3 times

mtmayer 1 year, 10 months ago

Selected Answer: A

The AD is on-premises... You need the connector.

upvoted 2 times

Mia2009687 1 year, 11 months ago

Selected Answer: D

<https://docs.aws.amazon.com/fsx/latest/WindowsGuide/aws-ad-integration-fsxW.html>

upvoted 2 times

kraken21 2 years, 2 months ago

Selected Answer: D

Other options are referring to IAM based control which is not possible. Existing AD should be used without IAM.

upvoted 3 times

Abhineet9148232 2 years, 2 months ago

Selected Answer: D

<https://aws.amazon.com/blogs/storage/using-amazon-fsx-for-windows-file-server-with-an-on-premises-active-directory/>

upvoted 3 times

somsundar 2 years, 2 months ago

Answer D. Amazon FSx does not support Active Directory Connector .

upvoted 3 times

JA2018 6 months, 3 weeks ago

<https://docs.aws.amazon.com/fsx/latest/WindowsGuide/aws-ad-integration-fsxW.html>
upvoted 1 times

Abhineet9148232 2 years, 3 months ago

Selected Answer: D

<https://docs.aws.amazon.com/fsx/latest/WindowsGuide/self-managed-AD.html>
upvoted 4 times

Yelizaveta 2 years, 3 months ago

Selected Answer: D

Note:
Amazon FSx does not support Active Directory Connector and Simple Active Directory.

<https://docs.aws.amazon.com/fsx/latest/WindowsGuide/aws-ad-integration-fsxW.html>
upvoted 4 times

aakashkumar1999 2 years, 4 months ago

Selected Answer: A

The answer will be AD connector so : A, it will create a proxy between your onpremises AD which you can use to restrict access
upvoted 2 times

JA2018 6 months, 3 weeks ago

from: <https://docs.aws.amazon.com/fsx/latest/WindowsGuide/aws-ad-integration-fsxW.html>

Amazon FSx does not support Active Directory Connector and Simple Active Directory.
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #261

Topic 1

A company recently announced the deployment of its retail website to a global audience. The website runs on multiple Amazon EC2 instances behind an Elastic Load Balancer. The instances run in an Auto Scaling group across multiple Availability Zones.

The company wants to provide its customers with different versions of content based on the devices that the customers use to access the website.

Which combination of actions should a solutions architect take to meet these requirements? (Choose two.)

- A. Configure Amazon CloudFront to cache multiple versions of the content. **Most Voted**
- B. Configure a host header in a Network Load Balancer to forward traffic to different instances.
- C. Configure a Lambda@Edge function to send specific objects to users based on the User-Agent header. **Most Voted**
- D. Configure AWS Global Accelerator. Forward requests to a Network Load Balancer (NLB). Configure the NLB to set up host-based routing to different EC2 instances.
- E. Configure AWS Global Accelerator. Forward requests to a Network Load Balancer (NLB). Configure the NLB to set up path-based routing to different EC2 instances.

Correct Answer: AC

Community vote distribution

AC (100%)

Comments

Parsons **Highly Voted** 1 year, 10 months ago

Selected Answer: AC

A, C is correct.

NLB listener rule only supports Protocol & Port (Not host/based routing like ALB) => D, E is incorrect.
NLB just works layer 4 (TCP/UDP) instead of Layer 7 (HTTP) => B is incorrect.

After eliminating, AC should be the answer.

upvoted 19 times

Guru4Cloud **Highly Voted** 1 year, 2 months ago

Selected Answer: AC

Selected Answer: AC

A. allows customers to receive the appropriate version of the content based on their location and device type.

C. By creating a Lambda@Edge, you can inspect the User-Agent header of incoming requests and determine the type of device being used. Based on this information, you can customize the response and send the appropriate version of the content to the user.

upvoted 11 times

Ruffyyit **Most Recent** 1 year ago

A C

Configure Amazon CloudFront to cache multiple versions of the content.

Configure a function to send specific objects to users based on the User-Agent header.

upvoted 3 times

sunhouse 1 year, 1 month ago

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/header-caching.html>

upvoted 3 times

rrbrish73 1 year, 2 months ago

<https://medium.com/swlh/serve-different-content-based-on-user-agent-in-aws-cloudfront-using-lambda-edge-28877294340b>

upvoted 2 times

cookieMr 1 year, 5 months ago

Selected Answer: AC

A. allows customers to receive the appropriate version of the content based on their location and device type.

C. By creating a Lambda@Edge, you can inspect the User-Agent header of incoming requests and determine the type of device being used. Based on this information, you can customize the response and send the appropriate version of the content to the user.

B. does not address the requirement of serving different content versions based on device types.

D. & E. do not address the device-specific content requirement.

Therefore, options A and C are the correct combination of actions to meet the requirement of providing different versions of content based on the devices that customers use to access the website.

upvoted 6 times

Yadav_Sanjay 1 year, 6 months ago

Selected Answer: AC

NLB does not supports routing

upvoted 3 times

omoakin 1 year, 6 months ago

A C

Configure Amazon CloudFront to cache multiple versions of the content.

Configure a function to send specific objects to users based on the User-Agent header.

upvoted 2 times

omoakin 1 year, 6 months ago

C

Configure a function to send specific objects to users based on the User-Agent header.

upvoted 1 times

GalileoEC2 1 year, 8 months ago

Using a Directory Connector to connect the on-premises Active Directory to AWS is one way to enable access to AWS resources, including Amazon FSx for Windows File Server. However, joining the Amazon FSx for Windows File Server file system to the on-premises Active Directory is a separate step that allows you to control access to the file shares using the same Active Directory groups that are used on-premises.

upvoted 1 times

LoXeras 1 year, 8 months ago

I guess this belongs to the question before #260

upvoted 3 times

wors 1 year, 10 months ago

So will this mean the entire architecture needs to move to lambda in order to leverage off lambda edge? This doesn't make sense as the question outlines the architecture already in ec2, asg and elb?

Just looking for clarification if I am missing something

upvoted 2 times

pentium75 11 months, 2 weeks ago

No, Lambda function will just do something like "if user-agent is 'iOS' then send content from elb-01 otherwise send content from elb-02".

upvoted 2 times

devonwho 1 year, 10 months ago

Selected Answer: AC

AC are the correct answers.

For C:

IMPROVED USER EXPERIENCE

Lambda@Edge can help improve your users' experience with your websites and web applications across the world, by letting you personalize content for them without sacrificing performance.

Real-time Image Transformation

You can customize your users' experience by transforming images on the fly based on the user characteristics. For example, you can resize images based on the viewer's device type—mobile, desktop, or tablet. You can also cache the transformed images at CloudFront Edge locations to further improve performance when delivering images.

<https://aws.amazon.com/lambda/edge/>

upvoted 3 times

mhmt4438 1 year, 10 months ago

Selected Answer: AC

Correct answer is A,C

upvoted 4 times

Aninina 1 year, 10 months ago

Selected Answer: AC

C. Configure a Lambda@Edge function to send specific objects to users based on the User-Agent header.

Lambda@Edge allows you to run a Lambda function in response to specific CloudFront events, such as a viewer request, an origin request, a response, or a viewer response.

upvoted 3 times

Morinator 1 year, 11 months ago

Selected Answer: AC

<https://www.examttopics.com/discussions/amazon/view/67881-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #262

Topic 1

A company plans to use Amazon ElastiCache for its multi-tier web application. A solutions architect creates a Cache VPC for the ElastiCache cluster and an App VPC for the application's Amazon EC2 instances. Both VPCs are in the us-east-1 Region.

The solutions architect must implement a solution to provide the application's EC2 instances with access to the ElastiCache cluster.

Which solution will meet these requirements MOST cost-effectively?

A. Create a peering connection between the VPCs. Add a route table entry for the peering connection in both VPCs. Configure an inbound rule for the ElastiCache cluster's security group to allow inbound connection from the application's security group. **Most Voted**

B. Create a Transit VPC. Update the VPC route tables in the Cache VPC and the App VPC to route traffic through the Transit VPC. Configure an inbound rule for the ElastiCache cluster's security group to allow inbound connection from the application's security group.

C. Create a peering connection between the VPCs. Add a route table entry for the peering connection in both VPCs. Configure an inbound rule for the peering connection's security group to allow inbound connection from the application's security group.

D. Create a Transit VPC. Update the VPC route tables in the Cache VPC and the App VPC to route traffic through the Transit VPC. Configure an inbound rule for the Transit VPC's security group to allow inbound connection from the application's security group.

Correct Answer: A

Community vote distribution

A (100%)

Comments

mhmt4438 **Highly Voted** 1 year, 4 months ago

Selected Answer: A

A. Create a peering connection between the VPCs. Add a route table entry for the peering connection in both VPCs. Configure an inbound rule for the ElastiCache cluster's security group to allow inbound connection from the application's security group.

Creating a peering connection between the VPCs allows the application's EC2 instances to communicate with the ElastiCache cluster directly and efficiently. This is the most cost-effective solution as it does not involve creating additional resources such as a Transit VPC, and it does not incur additional costs for traffic passing through the Transit VPC. Additionally, it is also more secure as it allows you to configure a more restrictive security group rule to allow inbound connection from only the application's security group.

upvoted 16 times

Ruffyt Most Recent 6 months, 3 weeks ago

A. Create a peering connection between the VPCs. Add a route table entry for the peering connection in both VPCs. Configure an inbound rule for the ElastiCache cluster's security group to allow inbound connection from the application's security group.

Creating a peering connection between the VPCs allows the application's EC2 instances to communicate with the ElastiCache cluster directly and efficiently. This is the most cost-effective solution as it does not involve creating additional resources such as a Transit VPC, and it does not incur additional costs for traffic passing through the Transit VPC. Additionally, it is also more secure as it allows you to configure a more restrictive security group rule to allow inbound connection from only the application's security group.

upvoted 3 times

TariqKipkemei 8 months, 2 weeks ago

Selected Answer: A

Create a peering connection between the VPCs. Add a route table entry for the peering connection in both VPCs. Configure an inbound rule for the ElastiCache cluster's security group to allow inbound connection from the application's security group.

upvoted 2 times

Guru4Cloud 9 months ago

Selected Answer: A

Create a VPC peering connection between the Cache VPC and App VPC. This allows private IP connectivity between the VPCs. Add route table entries in each VPC to route traffic destined to the other VPC via the peering connection. This enables network routing.

Configure security groups to allow inbound connections from the application instances to the ElastiCache cluster.

upvoted 2 times

cookieMr 11 months, 2 weeks ago

Selected Answer: A

Creating a peering connection between the VPCs is a cost-effective way to establish connectivity. By adding a route table entry for the peering connection in both VPCs, traffic can flow between them. Configuring an inbound rule in the ElastiCache cluster's security group allows inbound connections from the application's security group, enabling access to the ElastiCache cluster from the EC2 instances in the App VPC.

Option B suggests creating a Transit VPC, which adds unnecessary complexity and cost for this scenario.

Option C suggests configuring an inbound rule for the peering connection's security group, which is not necessary as the security group for the ElastiCache cluster should be used to control inbound connections.

Option D suggests configuring an inbound rule for the Transit VPC's security group, which is not needed in this case and adds unnecessary complexity.

Therefore, option A is the most cost-effective solution to provide the application's EC2 instances with access to the ElastiCache cluster.

upvoted 4 times

smartegnine 12 months ago

Selected Answer: A

A is correct,

1. VPC transit is used for more complex architecture and can do VPCs to VPCs connectivity. But for simple VPC 2 VPC can use peer connection.

2.To enable private IPv4 traffic between instances in peered VPCs, you must add a route to the route tables associated with the subnets for both instances.

So base on 1, B and D are out, base on 2 C is out

upvoted 2 times

wRhIH 1 year ago

Why not C ? any explanation?

upvoted 1 times

smartegnine 11 months, 3 weeks ago

Application read from ElasticCache, not viseversa, so inbound rule should be ElasticCach
upvoted 3 times

Cor5in 11 months, 2 weeks ago

Thank you Sir!
upvoted 1 times

smartegnine 12 months ago

To enable private IPv4 traffic between instances in peered VPCs, you must add a route to the route tables associated with the subnets for both instances.

<https://docs.aws.amazon.com/vpc/latest/peering/vpc-peering-routing.html>
upvoted 3 times

nder 1 year, 3 months ago

Selected Answer: A

Cost Effectively!
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #263

Topic 1


A company is building an application that consists of several microservices. The company has decided to use container technologies to deploy its software on AWS. The company needs a solution that minimizes the amount of ongoing effort for maintenance and scaling. The company cannot manage additional infrastructure.

Which combination of actions should a solutions architect take to meet these requirements? (Choose two.)

- A. Deploy an Amazon Elastic Container Service (Amazon ECS) cluster. **Most Voted**
- B. Deploy the Kubernetes control plane on Amazon EC2 instances that span multiple Availability Zones.
- C. Deploy an Amazon Elastic Container Service (Amazon ECS) service with an Amazon EC2 launch type. Specify a desired task number level of greater than or equal to 2.
- D. Deploy an Amazon Elastic Container Service (Amazon ECS) service with a Fargate launch type. Specify a desired task number level of greater than or equal to 2. **Most Voted**
- E. Deploy Kubernetes worker nodes on Amazon EC2 instances that span multiple Availability Zones. Create a deployment that specifies two or more replicas for each microservice.

Correct Answer: AD

Community vote distribution



Comments

cookieMr **Highly Voted** 1 year, 11 months ago

Selected Answer: AD

Options B and E suggest deploying the Kubernetes control plane and worker nodes on EC2 instances, which would require managing the infrastructure and add ongoing maintenance overhead, contrary to the requirement of minimizing effort.

Option C suggests using the Amazon EC2 launch type for ECS, which still requires managing EC2 instances and is not as cost-effective and scalable as using Fargate.

Therefore, the combination of deploying an Amazon ECS cluster and an ECS service with a Fargate launch type (options A and D) is the most suitable for minimizing maintenance and scaling effort without managing additional infrastructure.

upvoted 9 times

AlessandraSAA **Highly Voted** 2 years, 3 months ago

Selected Answer: AD

ECS has 2 launch type, EC2 (you maintain the infra) and Fargate (serverless). Since the question ask for no additional infra to manage it should be Fargate.

upvoted 5 times

TariqKipkemei **Most Recent** 1 year, 8 months ago

Selected Answer: AD

Company needs a solution that minimizes the amount of ongoing effort for maintenance and scaling = Serverless = ECS with Fargate.

upvoted 3 times

JA2018 6 months, 3 weeks ago

Rule of thumbs for AWS exams, always opt for AWS managed services

upvoted 1 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: AD

ECS allows deploying and managing containers without having to provision the underlying infrastructure. This minimizes maintenance effort.

Using Fargate launch type means ECS will handle provisioning and scaling the infrastructure automatically. This removes the management overhead for the company.

Running multiple tasks and specifying desired count ≥ 2 will provide high availability across Availability Zones.

Together, ECS plus Fargate provide a fully managed container platform. The company doesn't need to provision or manage servers.

upvoted 4 times

LoXeras 2 years, 2 months ago

Selected Answer: AD

AWS Fargate is server less solution to use on ECS:

https://docs.aws.amazon.com/AmazonECS/latest/developerguide/AWS_Fargate.html

upvoted 3 times

lambda15 2 years, 2 months ago

why is c is incorrect ?

upvoted 1 times

Julio98 2 years, 2 months ago

Because in the question says, "minimizes the amount of ongoing effort for maintenance and scaling", and EC2 instances you need effort to maintain the infrastructure unlike fargate that is serverless.

upvoted 3 times

WhericanIstart 2 years, 2 months ago

Selected Answer: AD

Amazon Fargate is a service that is fully manageable by Amazon; it offers provisioning, configuration and scaling feature. It is "serverless"..

upvoted 2 times

devonwho 2 years, 4 months ago

Selected Answer: AD

AWS Fargate is a technology that you can use with Amazon ECS to run containers without having to manage servers or clusters of Amazon EC2 instances. With Fargate, you no longer have to provision, configure, or scale clusters of virtual machines to run containers.

<https://docs.aws.amazon.com/AmazonECS/latest/userguide/what-is-fargate.html>

upvoted 4 times

Aninina 2 years, 4 months ago

A D is the correct answer

upvoted 2 times

mhmt4438 2 years, 4 months ago

Selected Answer: AD

A,D is correct answer

upvoted 3 times

AHUI 2 years, 4 months ago

AD:

<https://www.examttopics.com/discussions/amazon/view/60032-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 4 times

Morinator 2 years, 4 months ago

Selected Answer: AD

AD - EC2 out for this, cluster + fargate is the right answer

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #264


Topic 1

A company has a web application hosted over 10 Amazon EC2 instances with traffic directed by Amazon Route 53. The company occasionally experiences a timeout error when attempting to browse the application. The networking team finds that some DNS queries return IP addresses of unhealthy instances, resulting in the timeout error.

What should a solutions architect implement to overcome these timeout errors?

- A. Create a Route 53 simple routing policy record for each EC2 instance. Associate a health check with each record.
- B. Create a Route 53 failover routing policy record for each EC2 instance. Associate a health check with each record.
- C. Create an Amazon CloudFront distribution with EC2 instances as its origin. Associate a health check with the EC2 instances.

D. Create an Application Load Balancer (ALB) with a health check in front of the EC2 instances. Route to the ALB from Route 53. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****Guru4Cloud** **Highly Voted** 1 year, 9 months ago**Selected Answer: D**

ALB performs health checks on the EC2 instances, so it will only route traffic to healthy instances. This avoids the timeout errors. ALB provides load balancing across the instances, improving performance and availability. Route 53 routes to the ALB DNS name, so you don't have to manage records for each EC2 instance. This is a standard and robust architecture for public-facing web applications. The ALB acts as the entry point and handles health checks and scaling.

upvoted 16 times

jlteunissen **Highly Voted** 1 year, 9 months ago**Selected Answer: B**

It is not clear from the question whether the 10 EC2s are running within the same region. ALB can only direct traffic within region, while route 53 can route traffic to multiple locations, hence C and D are wrong.

upvoted 10 times

pentium75 1 year, 5 months ago

But B has one primary record and 9 failover records. A is correct, simple policy with health checks, that makes sure that only IPs of healthy instances are returned.

upvoted 3 times

JA2018 6 months, 3 weeks ago

err, can someone advise which is the correct answer?

upvoted 1 times

LeonSauveterre Most Recent 6 months, 2 weeks ago

Selected Answer: D

At first sight I thought option D is too complicated, but come to think of it, this should be the final answer. An ALB is specifically designed for this use case.

Option B is not suited for scenarios with multiple EC2 instances in the same environment because it works on a primary-secondary model, not load balancing across multiple instances.

upvoted 2 times

ChymKuBoy 11 months, 1 week ago

Selected Answer: D

D for sure

upvoted 2 times

Hrishi_707 1 year, 3 months ago

Those who are confused between A and D, A is wrong as you can't associate a health check with Simple routing policy record.

upvoted 2 times

MrPCarrot 1 year, 3 months ago

D is the best answer

upvoted 2 times

awsgeek75 1 year, 4 months ago

If you focus on the question, both A and D seems to be correct.

A is correct because simple routing policy for health check is doable BUT it is also wrong because we don't know how to determine the health of instance.

D is correct because " The company occasionally experiences a timeout error when attempting to browse the application" which suggest application is being accessed by a browser with means it's HTTP based and ALB is better for HTTP based healthchecks.

A web application timing out is not necessarily unhealthy instance, strictly speaking. It's just bad web application running on a healthy instance! So A may not be correct also.

upvoted 4 times

farnamjam 1 year, 5 months ago

Selected Answer: D

Although B can work as well, but it's nor a professional choice to associate the healthcheck with 10 EC2 instances, ALB is better option here.

A is incorect: Simple Routing Policies Can't be associated with Health Checks

C is incorrect: Cloudfront is for caching content which is irrelevant.

upvoted 4 times

awsgeek75 1 year, 5 months ago

Simple routing policies can be associated with Health Checks

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/dns-failover-simple-configs.html>

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: A

A meets the requirement ("overcome these timeout errors") without any other changes.

"If you configure health checking for all the records in a group of records that have the same name, the same type (such as A or AAAA), and the same routing policy (such as weighted or failover), Route 53 responds to DNS queries by choosing a healthy record and returning the applicable value from that record. (<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/health-checks-how-route-53-chooses-records.html>)

upvoted 3 times

pentium75 1 year, 5 months ago

"You can use Route 53 health checking to configure active-active and active-passive failover configurations. You configure active-active failover using any routing policy (or combination of routing policies) other than failover. ... Use this failover configuration when you want all of your resources to be available the majority of the time. When a resource becomes unavailable, Route 53 can detect that it's unhealthy and stop including it when responding to queries.

upvoted 2 times

pentium75 1 year, 5 months ago

In active-active failover, all the records that have the same name, the same type (such as A or AAAA), and the same routing policy (such as weighted or latency) are active unless Route 53 considers them unhealthy. Route 53 can respond to a DNS query using any healthy record." (<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/dns-failover-types.html#dns-failover-types-active-passive>)

upvoted 2 times

pentium75 1 year, 5 months ago

B - No, "you configure active-active failover using any routing policy (or combination of routing policies) OTHER THAN FAILOVER". With B, all traffic would go one primary instance unless it is unhealthy.

C - Not sure how to configure multiple EC2 instances as the origin without an LB. Even if that would be possible it would introduce more changes and complexity, which is not asked for.

D - Would work if all EC2 instances are in the same region, which we don't know. But it would also incur additional cost and potentially have other effects.

upvoted 2 times

Ruffyt 1 year, 6 months ago

B is wrong.

The DNS cache in clients could drive to timeouts. With ALB this issue won't happen since the DNS register will be the same and ALB will take care of unhealthy nodes.

upvoted 3 times

rlamberti 1 year, 7 months ago

Selected Answer: D

B is wrong.

The DNS cache in clients could drive to timeouts. With ALB this issue won't happen since the DNS register will be the same and ALB will take care of unhealthy nodes.

upvoted 3 times

daniel1 1 year, 7 months ago

Selected Answer: D

D. ****Application Load Balancer (ALB) with Health Checks, Routed via Route 53****:

- Creating an ALB in front of the EC2 instances and configuring health checks on the ALB will ensure that only healthy instances receive traffic. Route 53 can then direct traffic to the ALB, which in turn, routes traffic to healthy instances based on the health check results.

Among the provided options, the one that directly addresses the issue of routing traffic only to healthy instances is:

****D. Create an Application Load Balancer (ALB) with a health check in front of the EC2 instances. Route to the ALB from Route 53.****

upvoted 4 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: B

Clearly the question is all about Amazon Route 53 that has Failover routing policy that is used when you want to configure active-passive failover.

upvoted 1 times

slackbot 1 year, 9 months ago

i was looking at A, but indeed D is the best option, because the usually the TTL of the records is at least 60 seconds (nobody sets lower unless testing something ,because there is a charge per number of unique requests. ALB health check can be set as low as desired, which helps exclude the problematic ec2 faster than the DNS TTL expires

upvoted 3 times

cookieMr 1 year, 11 months ago

Selected Answer: D

By creating an ALB and configuring health checks, the architect ensures that only healthy instances receive traffic. The ALB periodically checks the health of the EC2 instances based on the configured health check settings.

Routing traffic to the ALB from Route 53 ensures that DNS queries return the IP address of the ALB instead of individual instances. This allows the ALB to distribute traffic only to healthy instances, avoiding timeouts caused by unhealthy instances.

A & B: While associating health checks with each record can help identify unhealthy instances, it does not provide automatic load balancing and distribution of traffic to healthy instances.

C: While CloudFront can improve performance and availability, it is primarily a CDN and may not directly address the issue of load balancing and distributing traffic to healthy instances.

Therefore, option D is the most appropriate solution to overcome the timeout errors by implementing an ALB with health checks and routing traffic through Route 53.

upvoted 5 times

joechen2023 1 year, 11 months ago

Selected Answer: C

I believe both C and D will work, but C seems less complex.

hopefully somebody here is more advanced(not an old student learning AWS like me) to explain why not C.

upvoted 3 times

Abrar2022 2 years ago

Selected Answer: D

Option D allows for the creation of an Application Load Balancer which can detect unhealthy instances and redirect traffic away from them.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #265

Topic 1

A solutions architect needs to design a highly available application consisting of web, application, and database tiers. HTTPS content delivery should be as close to the edge as possible, with the least delivery time.

Which solution meets these requirements and is MOST secure?

- A. Configure a public Application Load Balancer (ALB) with multiple redundant Amazon EC2 instances in public subnets. Configure Amazon CloudFront to deliver HTTPS content using the public ALB as the origin.
- B. Configure a public Application Load Balancer with multiple redundant Amazon EC2 instances in private subnets. Configure Amazon CloudFront to deliver HTTPS content using the EC2 instances as the origin.
- C. Configure a public Application Load Balancer (ALB) with multiple redundant Amazon EC2 instances in private subnets. Configure Amazon CloudFront to deliver HTTPS content using the public ALB as the origin. **Most Voted**
- D. Configure a public Application Load Balancer with multiple redundant Amazon EC2 instances in public subnets. Configure Amazon CloudFront to deliver HTTPS content using the EC2 instances as the origin.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Aninina** **Highly Voted** 2 years, 4 months ago

C. Configure a public Application Load Balancer (ALB) with multiple redundant Amazon EC2 instances in private subnets. Configure Amazon CloudFront to deliver HTTPS content using the public ALB as the origin.

This solution meets the requirements for a highly available application with web, application, and database tiers, as well as providing edge-based content delivery. Additionally, it maximizes security by having the ALB in a private subnet, which limits direct access to the web servers, while still being able to serve traffic over the Internet via the public ALB. This will ensure that the web servers are not exposed to the public Internet, which reduces the attack surface and provides a secure way to access the application.

upvoted 18 times

JA2018 6 months, 3 weeks ago

err, how do we know there's already an ALB in the public subnet(s)?

upvoted 1 times

JA2018 6 months, 3 weeks ago

This is *NOT* stated in the sample question listed here.

upvoted 1 times

JA2018 6 months, 3 weeks ago

sorry missed out on this "Configure Amazon CloudFront to deliver HTTPS content using the public ALB as the origin."

upvoted 1 times

cookieMr **Highly Voted** 1 year, 11 months ago

Selected Answer: C

A. exposes the EC2 instances directly to the public internet, which may compromise security.

B. lacks a load balancer in the public subnet, which is required for efficient load distribution and high availability.

D. provides load balancing and HTTPS content delivery, it exposes the EC2 instances directly to the public internet, which may pose security risks.

C. provides high availability, secure access through private subnets, and optimized HTTPS content delivery using CloudFront with a public ALB as the origin.

upvoted 10 times

meowruki **Most Recent** 1 year, 6 months ago

C. Configure a public Application Load Balancer (ALB) with multiple redundant Amazon EC2 instances in private subnets. Configure Amazon CloudFront to deliver HTTPS content using the public ALB as the origin.

Here's the reasoning:

Public ALB in Private Subnets: Placing the ALB in private subnets enhances security by preventing direct access from the internet. The ALB in private subnets can communicate with the application instances in the same private subnets.

CloudFront with ALB as Origin: Configuring CloudFront to deliver HTTPS content using the public ALB as the origin allows for content to be cached and distributed globally, reducing latency for end users.

upvoted 2 times

Ruffyyit 1 year, 6 months ago

C. Configure a public Application Load Balancer (ALB) with multiple redundant Amazon EC2 instances in private subnets. Configure Amazon CloudFront to deliver HTTPS content using the public ALB as the origin.

This solution meets the requirements for a highly available application with web, application, and database tiers, as well as providing edge-based content delivery. Additionally, it maximizes security by having the ALB in a private subnet, which limits direct access to the web servers, while still being able to serve traffic over the Internet via the public ALB. This will ensure that the web servers are not exposed to the public Internet, which reduces the attack surface and provides a secure way to access the application.

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

Keyword: Instances in private, ALB in public, point cloudfront to the public ALB

upvoted 2 times

mhmt4438 2 years, 4 months ago

Selected Answer: C

Answer is C

upvoted 4 times

AHUI 2 years, 4 months ago

ans: C

<https://www.examttopics.com/discussions/amazon/view/46401-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Morinator 2 years, 4 months ago

Selected Answer: C

Instances in private, ALB in public, point cloudfront to the public ALB
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #266

Topic 1

A company has a popular gaming platform running on AWS. The application is sensitive to latency because latency can impact the user experience and introduce unfair advantages to some players. The application is deployed in every AWS Region. It runs on Amazon EC2 instances that are part of Auto Scaling groups configured behind Application Load Balancers (ALBs). A solutions architect needs to implement a mechanism to monitor the health of the application and redirect traffic to healthy endpoints.

Which solution meets these requirements?

- A. Configure an accelerator in AWS Global Accelerator. Add a listener for the port that the application listens on, and attach it to a Regional endpoint in each Region. Add the ALB as the endpoint. **Most Voted**
- B. Create an Amazon CloudFront distribution and specify the ALB as the origin server. Configure the cache behavior to use origin cache headers. Use AWS Lambda functions to optimize the traffic.
- C. Create an Amazon CloudFront distribution and specify Amazon S3 as the origin server. Configure the cache behavior to use origin cache headers. Use AWS Lambda functions to optimize the traffic.
- D. Configure an Amazon DynamoDB database to serve as the data store for the application. Create a DynamoDB Accelerator (DAX) cluster to act as the in-memory cache for DynamoDB hosting the application data.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Aninina **Highly Voted** 1 year, 10 months ago

Selected Answer: A

A. Configure an accelerator in AWS Global Accelerator. Add a listener for the port that the application listens on, and attach it to a Regional endpoint in each Region. Add the ALB as the endpoint.

AWS Global Accelerator directs traffic to the optimal healthy endpoint based on health checks, it can also route traffic to the closest healthy endpoint based on geographic location of the client. By configuring an accelerator and attaching it to a Regional endpoint in each Region, and adding the ALB as the endpoint, the solution will redirect traffic to healthy endpoints, improving the user experience by reducing latency and ensuring that the application is running optimally. This solution will ensure that traffic is directed to the closest healthy endpoint and will help to improve the overall user experience.

invited 22 times

Updated 22 minutes ago

michellemeloc Highly Voted 1 year, 6 months ago

Selected Answer: A

Delivery gaming content --> AWS GLOBAL ACCELERATOR

upvoted 15 times

pentium75 Most Recent 11 months, 2 weeks ago

Selected Answer: A

Would have selected A just because B, C and D don't make any sense or have nothing to do with the requirements. But now learned that Global Accelerator checks health of resources BEHIND ALB/NLB, so it meets the requirements.

upvoted 4 times

mwwt2022 11 months, 2 weeks ago

gaming platform -> Can't be CloudFront. Probably go for global accelerator

upvoted 5 times

Ruffyit 1 year ago

A. Configure an accelerator in AWS Global Accelerator. Add a listener for the port that the application listens on, and attach it to a Regional endpoint in each Region. Add the ALB as the endpoint.

AWS Global Accelerator directs traffic to the optimal healthy endpoint based on health checks, it can also route traffic to the closest healthy endpoint based on geographic location of the client. By configuring an accelerator and attaching it to a Regional endpoint in each Region, and adding the ALB as the endpoint, the solution will redirect traffic to healthy endpoints, improving the user experience by reducing latency and ensuring that the application is running optimally. This solution will ensure that traffic is directed to the closest healthy endpoint and will help to improve the overall user experience.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

A. Configure an accelerator in AWS Global Accelerator. Add a listener for the port that the application listens on, and attach it to a Regional endpoint in each Region. Add the ALB as the endpoint

upvoted 2 times

bjexamprep 1 year, 4 months ago

Is any answer relevant to the question?

upvoted 3 times

cookieMr 1 year, 5 months ago

Selected Answer: A

B. While CloudFront can help with caching and content delivery, it does not provide the mechanism to monitor the health of the application or perform traffic redirection based on health checks.

C. This configuration is suitable for static content delivery but does not address the health monitoring and traffic redirection requirements of the application.

D. While this can enhance performance, it does not monitor the health of the application or redirect traffic based on health checks.

Therefore, option A is the most suitable solution as it leverages AWS Global Accelerator to monitor application health, route traffic to healthy endpoints, and optimize the user experience while addressing latency concerns.

upvoted 6 times

antropaws 1 year, 6 months ago

Selected Answer: A

Agree with A

upvoted 2 times

Bhrino 1 year, 9 months ago

Selected Answer: A

Global accelerators can be used for non http cases such as UDP, tcp , gaming , or voip

upvoted 9 times

pentium75 11 months, 2 weeks ago

Though we seem to have http/https here, otherwise they could not use ALBs

upvoted 2 times

mhmt4438 1 year, 10 months ago

Selected Answer: A

Correct answer is A

upvoted 2 times

AHUI 1 year, 10 months ago

A:

<https://www.examtopycs.com/discussions/amazon/view/46403-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

alanp 1 year, 10 months ago

A. When you have an Application Load Balancer or Network Load Balancer that includes multiple target groups, Global Accelerator considers the load balancer endpoint to be healthy only if each target group behind the load balancer has at least one healthy target. If any single target group for the load balancer has only unhealthy targets, Global Accelerator considers the endpoint to be unhealthy.

<https://docs.aws.amazon.com/global-accelerator/latest/dg/about-endpoint-groups-health-check-options.html>

upvoted 9 times

Morinator 1 year, 11 months ago

Selected Answer: A

<https://docs.aws.amazon.com/global-accelerator/latest/dg/about-endpoint-groups-health-check-options.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #267

Topic 1

A company has one million users that use its mobile app. The company must analyze the data usage in near-real time. The company also must encrypt the data in near-real time and must store the data in a centralized location in Apache Parquet format for further processing.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an Amazon Kinesis data stream to store the data in Amazon S3. Create an Amazon Kinesis Data Analytics application to analyze the data. Invoke an AWS Lambda function to send the data to the Kinesis Data Analytics application.
- B. Create an Amazon Kinesis data stream to store the data in Amazon S3. Create an Amazon EMR cluster to analyze the data. Invoke an AWS Lambda function to send the data to the EMR cluster.
- C. Create an Amazon Kinesis Data Firehose delivery stream to store the data in Amazon S3. Create an Amazon EMR cluster to analyze the data.
- D. Create an Amazon Kinesis Data Firehose delivery stream to store the data in Amazon S3. Create an Amazon Kinesis Data Analytics application to analyze the data. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

mhmt4438 **Highly Voted** 1 year, 10 months ago

Selected Answer: D

D. Create an Amazon Kinesis Data Firehose delivery stream to store the data in Amazon S3. Create an Amazon Kinesis Data Analytics application to analyze the data.

This solution will meet the requirements with the least operational overhead as it uses Amazon Kinesis Data Firehose, which is a fully managed service that can automatically handle the data collection, data transformation, encryption, and data storage in near-real time. Kinesis Data Firehose can automatically store the data in Amazon S3 in Apache Parquet format for further processing. Additionally, it allows you to create an Amazon Kinesis Data Analytics application to analyze the data in near real-time, with no need to manage any infrastructure or invoke any Lambda function. This way you can process a large amount of data with the least operational overhead.

upvoted 56 times

antonyne 1 year, 6 months ago

anuropaws 1 year, 6 months ago

<https://aws.amazon.com/blogs/big-data/analyzing-apache-parquet-optimized-data-using-amazon-kinesis-data-firehose-amazon-athena-and-amazon-redshift/>
upvoted 2 times

WherecanIstart 1 year, 8 months ago

Thanks for the explanation!
upvoted 2 times

LuckyAro 1 year, 10 months ago

Apache Parquet format processing was not mentioned in the answer options. Strange.
upvoted 7 times

jainparag1 1 year, 10 months ago

Nicely explained. Thanks.
upvoted 3 times

cookieMr Highly Voted 1 year, 5 months ago

Selected Answer: D

A. requires invoking an Lambda to send the data to the analytics application. This introduces additional operational overhead and complexity.

B. While EMR is a powerful tool for big data processing, it requires more operational management and configuration compared to Kinesis Data Analytics.

C. introduces unnecessary complexity by involving EMR for data analysis when Kinesis Data Analytics can perform the analysis in a more streamlined and automated manner.

Therefore, option D is the most suitable solution as it leverages Kinesis Data Firehose for data ingestion, stores the data in S3, and utilizes Kinesis Data Analytics for near-real-time analysis, providing a low operational overhead solution for data usage analysis and encryption.

upvoted 8 times

farnamjam Most Recent 10 months, 2 weeks ago

Selected Answer: D

A and B are out.
Kinesis Data Streams cannot directly send data to S3 by itself
upvoted 3 times

Ruffyt 1 year ago

D. Create an Amazon Kinesis Data Firehose delivery stream to store the data in Amazon S3. Create an Amazon Kinesis Data Analytics application to analyze the data.

This solution will meet the requirements with the least operational overhead as it uses Amazon Kinesis Data Firehose, which is a fully managed service that can automatically handle the data collection, data transformation, encryption, and data storage in near-real time. Kinesis Data Firehose can automatically store the data in Amazon S3 in Apache Parquet format for further processing. Additionally, it allows you to create an Amazon Kinesis Data Analytics application to analyze the data in near real-time, with no need to manage any infrastructure or invoke any Lambda function. This way you can process a large amount of data with the least operational overhead.

upvoted 3 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: D

Create an Amazon Kinesis Data Firehose delivery stream to store the data in Amazon S3. Create an Amazon Kinesis Data Analytics application to analyze the data
upvoted 2 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: D

D. Create an Amazon Kinesis Data Firehose delivery stream to store the data in Amazon S3. Create an Amazon Kinesis Data Analytics application to analyze the data

upvoted 2 times

AHUI 1 year, 10 months ago

D:

<https://www.examttopics.com/discussions/amazon/view/82022-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 1 times

Aninina 1 year, 10 months ago

Selected Answer: D

D. Create an Amazon Kinesis Data Firehose delivery stream to store the data in Amazon S3. Create an Amazon Kinesis Data Analytics application to analyze the data.

Amazon Kinesis Data Firehose can automatically encrypt and store the data in Amazon S3 in Apache Parquet format for further processing, which reduces the operational overhead. It also allows for near-real-time data analysis using Kinesis Data Analytics, which is a fully managed service that makes it easy to analyze streaming data using SQL. This solution eliminates the need for setting up and maintaining an EMR cluster, which would require more operational overhead.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #268

Topic 1

A gaming company has a web application that displays scores. The application runs on Amazon EC2 instances behind an Application Load Balancer. The application stores data in an Amazon RDS for MySQL database. Users are starting to experience long delays and interruptions that are caused by database read performance. The company wants to improve the user experience while minimizing changes to the application's architecture.

What should a solutions architect do to meet these requirements?

- A. Use Amazon ElastiCache in front of the database. **Most Voted**
- B. Use RDS Proxy between the application and the database.
- C. Migrate the application from EC2 instances to AWS Lambda.
- D. Migrate the database from Amazon RDS for MySQL to Amazon DynamoDB.

Correct Answer: A*Community vote distribution***Comments****Steve_4542636** **Highly Voted** 2 years, 3 months ago**Selected Answer: A**

Rds proxy is for too many connections, not for performance
upvoted 35 times

Mkenya08 1 year, 4 months ago

ElastiCache stores data in memory, which means it may not always have the most up-to-date information. This might be acceptable for certain use cases where slightly stale data is acceptable, but for applications like gaming scores, real-time accuracy is often crucial.

upvoted 5 times

Maru86 1 year, 3 months ago

"Data in the cache is never stale.

Because the data in the cache is updated every time it's written to the database, the data in the cache is always current and updated whenever data is written to the database." <https://docs.aws.amazon.com/AmazonElastiCache/latest/mem->