

ug/Strategies.html#Strategies.WriteThrough
upvoted 4 times

vipyodha 1 year, 11 months ago

to use elasticache , you need to perform heavy code change ,and also elasticache do chaching that can improve read performance but will not provide scalability
upvoted 9 times

pentium75 1 year, 5 months ago

We should "minimize", not "avoid", code changes.
upvoted 3 times

chickenmf 1 year, 2 months ago

minimize *architectural changes, NOT code changes
upvoted 2 times

JackyCCK 1 year, 2 months ago

RDS Proxy: It helps applications improve scalability and performance by managing database connections and pooling, which can significantly reduce the load on the database
upvoted 3 times

Yadav_Sanjay 2 years ago

Can't use cache as score gates updated. If data would have been static then definitely can go with A. But here score is dynamic...
upvoted 10 times

rfelipem 2 years ago

Users are starting to experience long delays and interruptions caused by the "read performance" of the database... While the score is dynamic, there is also read activity in the DB that is causing the delays and outages and this can be improved with Elastic Cache.
upvoted 5 times

kraken21 **Highly Voted** 2 years, 2 months ago

Selected Answer: B

RDX proxy will : "improve the user experience while minimizing changes".
upvoted 26 times

pentium75 1 year, 5 months ago

.. but not address issues with "database read performance".
upvoted 3 times

thiahthura **Most Recent** 6 months ago

Selected Answer: A

users are experiencing long delays. This is the priority issue , so we need to use cache service urgently to solve this issues. If we need performance , we can use proxy. But in this case , latency is our priority. So Answer is A.
upvoted 1 times

LeonSauveterre 6 months, 2 weeks ago

Selected Answer: A

A: Correct. Not suitable for write-heavy applications though.
B: Reduces pressure on the database by reusing connections, but cannot address the slow-read problem.
C: Why? Who would do this?
D: It completely changes the architecture.
upvoted 1 times

Mish 7 months ago

Selected Answer: A

Use Amazon ElastiCache in front of the database.

upvoted 2 times

pujithacg8 10 months, 1 week ago

A. Use Amazon ElastiCache in front of the database.

Explanation:

Amazon ElastiCache is a managed caching service that can be placed in front of the database to cache frequently accessed data. By caching the most common queries, it can significantly reduce the load on the database, leading to faster response times and an improved user experience.

This solution minimizes changes to the existing architecture because it doesn't require modifications to the application's core logic or database schema. Instead, it optimizes performance by reducing the need for repetitive database reads.

ElastiCache supports both Redis and Memcached, which are widely used for caching in web applications.

upvoted 2 times

ChymKuBoy 11 months, 1 week ago

Selected Answer: A

A for sure

upvoted 2 times

Mayank0502 11 months, 2 weeks ago

Selected Answer: B

answer should be B

upvoted 1 times

Manjeet_Kumar 11 months, 2 weeks ago

Selected Answer: A

Not B because this can help with connection management and improve scalability and availability but won't directly address read performance issues caused by high read traffic.

upvoted 2 times

MomenAWS 12 months ago

Selected Answer: A

Another vague question from AWS: I would prefer A over B

upvoted 2 times

[Removed] 12 months ago

A

Caching frequently accessed data: ElastiCache can be used to reduce the load on your database by caching frequently accessed data. This can improve application performance and reduce the number of read queries to your database.

Real-time applications: If your application requires real-time data processing, ElastiCache can help. Since it provides sub-millisecond response times, it can be used to power applications like gaming leaderboards, chat applications, and real-time analytics.

upvoted 2 times

rohitph 1 year ago

Selected Answer: B

ask is - minimum changes.

RDS Proxy is a feature of Amazon RDS.

This would be easiest option to try.

upvoted 1 times

MandAsh 1 year ago

Selected Answer: A

Elasticache - reduce load due to read operations

RDS proxy - reduce load due to lot of connections

Here problem is read operations thus A is solution

upvoted 3 times

hb0011 1 year ago

Selected Answer: A

I vote for Elasticache

upvoted 2 times

rjjkc 1 year ago

Selected Answer: B

B - "Use RDS Proxy between the application and the database."

B because it improves the user experience while minimizing the changes.

If A is used, you have to modify your application to get the data from the cache first, if it is not there, then get from db; and also to invalidate the cache if there is a db update.

upvoted 1 times

osmk 1 year, 2 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AmazonElastiCache/latest/APIReference/Welcome.html>

upvoted 2 times

Uzbekistan 1 year, 2 months ago

Selected Answer: B

RDS Proxy: It helps applications improve scalability and performance by managing database connections and pooling, which can significantly reduce the load on the database. By using RDS Proxy, you can offload connection management tasks from your application and optimize database access, thereby improving read performance without making significant changes to the application's architecture.

Minimizing Changes: Implementing RDS Proxy does not require architectural changes to the application. It works as a proxy layer between the application and the RDS database, so the application code remains unchanged. This minimizes the effort and risk associated with making modifications to the application.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #269

Topic 1

An ecommerce company has noticed performance degradation of its Amazon RDS based web application. The performance degradation is attributed to an increase in the number of read-only SQL queries triggered by business analysts. A solutions architect needs to solve the problem with minimal changes to the existing web application.

What should the solutions architect recommend?

- A. Export the data to Amazon DynamoDB and have the business analysts run their queries.
- B. Load the data into Amazon ElastiCache and have the business analysts run their queries.
- C. Create a read replica of the primary database and have the business analysts run their queries. **Most Voted**
- D. Copy the data into an Amazon Redshift cluster and have the business analysts run their queries.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Aninina** **Highly Voted** 1 year, 10 months ago**Selected Answer: C**

C. Create a read replica of the primary database and have the business analysts run their queries.

Creating a read replica of the primary RDS database will offload the read-only SQL queries from the primary database, which will help to improve the performance of the web application. Read replicas are exact copies of the primary database that can be used to handle read-only traffic, which will reduce the load on the primary database and improve the performance of the web application. This solution can be implemented with minimal changes to the existing web application, as the business analysts can continue to run their queries on the read replica without modifying the code.

upvoted 6 times

awsgeek75 **Most Recent** 10 months, 3 weeks ago**Selected Answer: C**

". A solutions architect needs to solve the problem with minimal changes to the existing web application."

ABD all require major changes to the application.

A: DynamoDB is NoSQL so big change

B: ElastiCache is a caching layer which require code change normally significant code change

D: Redshift is analytics so not a solution

upvoted 4 times

Ruffyt 1 year ago

. While DynamoDB is a scalable NoSQL database, it requires changes to the application's data model and query patterns.

B. ElastiCache is an in-memory data store that can improve query performance, but it is primarily used for caching rather than running complex queries.

D. Redshift is a powerful data warehousing solution, but migrating the data and adapting the queries to Redshift's columnar architecture would require significant changes to the application and query logic.

Therefore, option C is the most appropriate recommendation as it leverages read replicas in RDS to offload read-only query traffic from the primary database, allowing the business analysts to run their queries without impacting the performance of the web application. It provides a scalable and efficient solution with minimal changes to the existing web application.

upvoted 2 times

nileeka97 1 year, 2 months ago

Selected Answer: C

C. Create a read replica of the primary database and have the business analysts run their queries

upvoted 2 times

cookieMr 1 year, 5 months ago

Selected Answer: C

A. While DynamoDB is a scalable NoSQL database, it requires changes to the application's data model and query patterns.

B. ElastiCache is an in-memory data store that can improve query performance, but it is primarily used for caching rather than running complex queries.

D. Redshift is a powerful data warehousing solution, but migrating the data and adapting the queries to Redshift's columnar architecture would require significant changes to the application and query logic.

Therefore, option C is the most appropriate recommendation as it leverages read replicas in RDS to offload read-only query traffic from the primary database, allowing the business analysts to run their queries without impacting the performance of the web application. It provides a scalable and efficient solution with minimal changes to the existing web application.

upvoted 3 times

antropaws 1 year, 6 months ago

Selected Answer: C

C, no doubt.

upvoted 3 times

mhmt4438 1 year, 10 months ago

Selected Answer: C

C is correct answer

upvoted 3 times

bamishr 1 year, 11 months ago

Selected Answer: C

Create a read replica of the primary database and have the business analysts run their queries.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #270

Topic 1

A company is using a centralized AWS account to store log data in various Amazon S3 buckets. A solutions architect needs to ensure that the data is encrypted at rest before the data is uploaded to the S3 buckets. The data also must be encrypted in transit.

Which solution meets these requirements?

- A. Use client-side encryption to encrypt the data that is being uploaded to the S3 buckets. **Most Voted**
- B. Use server-side encryption to encrypt the data that is being uploaded to the S3 buckets.
- C. Create bucket policies that require the use of server-side encryption with S3 managed encryption keys (SSE-S3) for S3 uploads.
- D. Enable the security option to encrypt the S3 buckets through the use of a default AWS Key Management Service (AWS KMS) key.

Correct Answer: A*Community vote distribution*

A (94%)

C (6%)

Comments**techhb** **Highly Voted** 2 years, 4 months ago**Selected Answer: A**

here keyword is "before" "the data is encrypted at rest before the data is uploaded to the S3 buckets."

upvoted 29 times

mhmt4438 **Highly Voted** 2 years, 4 months ago**Selected Answer: A**

<https://www.examtopycs.com/discussions/amazon/view/53840-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 5 times

Dantecito **Most Recent** 4 months ago**Selected Answer: A**

For those who are confused with the data also must be encrypted in transit, Amazon S3 TLS (Transport Layer Security) is used by default and you can't disable it.

default and you can't disable it.

upvoted 2 times

LeonSauveterre 6 months, 2 weeks ago

Selected Answer: A

After a little googling, for those who are confused by "at rest":

1. *Encryption at Rest* refers to the encryption applied to the stored data. Encryption may be implemented at the source, where data is generated and stored at the origin.
2. *Encryption in Transit* refers to encrypting data that is transferred between two nodes of the network.
3. *End-to-End Encryption* refers to the combination of the encryption at rest and encryption in transit.

upvoted 3 times

babayomi 7 months, 3 weeks ago

A
I believe the question is crafted to cause some confusion. At the same time it is simple to answer, since client side encryption answers the requirements.

upvoted 2 times

reviewmine 1 year, 3 months ago

Selected Answer: A

Answer is A. Encrypt it first before uploading to S3.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: C

I think the many votes for A are caused by misunderstanding the wording as
"Ensure that
the data is encrypted at rest before the data is uploaded"

But that doesn't make sense, it means

"Ensure that the data is encrypted at rest
before the data is uploaded"

So, before you allow people to upload data, make sure that it gets encrypted.

upvoted 3 times

pentium75 1 year, 5 months ago

On second thought, C would not enforce encryption in transit. Thus must be A indeed.

upvoted 3 times

awsgeek75 1 year, 4 months ago

For a moment I bought into your reasoning for C assuming that maybe the question is missing some grammar construct but realised that C does not really solve the encryption in transit issue like I originally thought. BUT good work!

upvoted 2 times

Cyberkayu 1 year, 5 months ago

BCD, data not yet encrypted before landing on S3 bucket

upvoted 3 times

palthainon 1 year, 7 months ago

Selected Answer: C

HTTPs would encrypt in transe, SSE3 managed keys fulfills requirement for at rest. This is an aws exam, not a best practices exam.

upvoted 1 times

awsgeek75 1 year, 4 months ago

No. HTTPS is not enough for encryption in transit when it comes to S3.
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingClientSideEncryption.html>

"Client-side encryption is the act of encrypting your data locally to help ensure its security in transit and at rest. "
upvoted 4 times

petertang224 1 year, 7 months ago

Its_SaKar
upvoted 1 times

prabhjot 1 year, 8 months ago

Ans is B - Server-Side Encryption (SSE): ensure data is encrypted at rest and also Encryption in Transit: When you upload data to Amazon S3 using standard HTTPS requests.
upvoted 4 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: A

Use client-side encryption to encrypt the data that is being uploaded to the S3 buckets
upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

A. Use client-side encryption to encrypt the data that is being uploaded to the S3 buckets.
upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

A. Use client-side encryption to encrypt the data that is being uploaded to the S3 buckets.
upvoted 2 times

Abobaloyi 1 year, 11 months ago

Selected Answer: A

data must be encrypted before uploaded , which means the client need to do it before uploading the data to S3
upvoted 4 times

datz 2 years, 2 months ago

Selected Answer: A

A, would meet requirements.
upvoted 2 times

nder 2 years, 3 months ago

Selected Answer: A

Because the data must be encrypted while in transit
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #271

Topic 1

A solutions architect observes that a nightly batch processing job is automatically scaled up for 1 hour before the desired Amazon EC2 capacity is reached. The peak capacity is the 'same every night and the batch jobs always start at 1 AM. The solutions architect needs to find a cost-effective solution that will allow for the desired EC2 capacity to be reached quickly and allow the Auto Scaling group to scale down after the batch jobs are complete.

What should the solutions architect do to meet these requirements?

- A. Increase the minimum capacity for the Auto Scaling group.
- B. Increase the maximum capacity for the Auto Scaling group.
- C. Configure scheduled scaling to scale up to the desired compute level. **Most Voted**
- D. Change the scaling policy to add more EC2 instances during each scaling operation.

Correct Answer: C

Community vote distribution

C (100%)

Comments

ManOnTheMoon **Highly Voted** 1 year, 4 months ago

GOOD LUCK EVERYONE :) YOU CAN DO THIS
upvoted 43 times

david76x **Highly Voted** 1 year, 4 months ago

Selected Answer: C

C is correct. Goodluck everybody!
upvoted 16 times

Guru4Cloud **Most Recent** 9 months ago

Selected Answer: C

Configuring scheduled scaling actions allows the Auto Scaling group to scale up to the desired capacity at a scheduled time (1 AM in this case) when the batch jobs start. This ensures the desired compute capacity is reached immediately.

The Auto Scaling group can then scale down based on metrics after the batch jobs complete.

upvoted 9 times

hsinchang 10 months, 2 weeks ago

Selected Answer: C

The time is given, use scheduled for optimal cost

upvoted 4 times

qacollin 1 year, 1 month ago

just scheduled my exam :)

upvoted 7 times

awscerts023 1 year, 3 months ago

Reached here ! Did anyone schedule the real exam now ? How was it ?

upvoted 6 times

pal40sg 1 year, 3 months ago

Thanks to everyone who contributed with answers :)

upvoted 5 times

ProfXsamson 1 year, 4 months ago

Selected Answer: C

C. I'm here at the end, leaving this here for posterity sake 02/01/2023.

upvoted 5 times

dedline 1 year, 4 months ago

GL ALL!

upvoted 5 times

mhmt4438 1 year, 4 months ago

Selected Answer: C

<https://www.examttopics.com/discussions/amazon/view/27868-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Aninina 1 year, 4 months ago

Selected Answer: C

C. Configure scheduled scaling to scale up to the desired compute level.

By configuring scheduled scaling, the solutions architect can set the Auto Scaling group to automatically scale up to the desired compute level at a specific time (1AM) when the batch job starts and then automatically scale down after the job is complete. This will allow the desired EC2 capacity to be reached quickly and also help in reducing the cost.

upvoted 5 times

bamishr 1 year, 4 months ago

Selected Answer: C

Configure scheduled scaling to scale up to the desired compute level.

upvoted 2 times

Morinator 1 year, 4 months ago

Selected Answer: C

predictable = schedule scaling

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #272

Topic 1

A company serves a dynamic website from a fleet of Amazon EC2 instances behind an Application Load Balancer (ALB). The website needs to support multiple languages to serve customers around the world. The website's architecture is running in the us-west-1 Region and is exhibiting high request latency for users that are located in other parts of the world.

The website needs to serve requests quickly and efficiently regardless of a user's location. However, the company does not want to recreate the existing architecture across multiple Regions.

What should a solutions architect do to meet these requirements?

- A. Replace the existing architecture with a website that is served from an Amazon S3 bucket. Configure an Amazon CloudFront distribution with the S3 bucket as the origin. Set the cache behavior settings to cache based on the Accept-Language request header.
- B. Configure an Amazon CloudFront distribution with the ALB as the origin. Set the cache behavior settings to cache based on the Accept-Language request header. **Most Voted**
- C. Create an Amazon API Gateway API that is integrated with the ALB. Configure the API to use the HTTP integration type. Set up an API Gateway stage to enable the API cache based on the Accept-Language request header.
- D. Launch an EC2 instance in each additional Region and configure NGINX to act as a cache server for that Region. Put all the EC2 instances and the ALB behind an Amazon Route 53 record set with a geolocation routing policy.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Yechi **Highly Voted** 2 years, 3 months ago

Selected Answer: B

Configuring caching based on the language of the viewer

If you want CloudFront to cache different versions of your objects based on the language specified in the request, configure CloudFront to forward the Accept-Language header to your origin.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/header-caching.html>

upvoted 13 times

[Removed] **Most Recent** 12 months ago

B

If you want CloudFront to cache different versions of your objects based on the language specified in the request, configure CloudFront to forward the Accept-Language header to your origin.

If you want CloudFront to cache different versions of your objects based on the country that the request came from, configure CloudFront to forward the CloudFront-Viewer-Country header to your origin. CloudFront automatically converts the IP address that the request came from into a two-letter country code. For an easy-to-use list of country codes, sortable by code and by country name, see the Wikipedia entry ISO 3166-1 alpha-2.

upvoted 3 times

Trains 1 year, 6 months ago

Isn't CloudFront for static websites though? Question specifically states the content is dynamic

upvoted 2 times

Cloud_A 1 year, 4 months ago

Cloudfront serves for both static and dynamic. If it was just static, then you can consider AWS S3.

upvoted 5 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

By caching content based on the Accept-Language request header, CloudFront can serve the appropriate version of the website to users based on their language preferences. This solution allows the company to improve the website's performance for users around the world without having to recreate the existing architecture in multiple Regions.

upvoted 4 times

A1975 1 year, 10 months ago

Selected Answer: B

CloudFront allows you to customize cache behavior based on various request headers. By setting the cache behavior to cache based on the Accept-Language request header, CloudFront can store and serve language-specific versions of the website content, reducing the need to repeatedly fetch the content from the ALB for users with the same language preference.

upvoted 3 times

kraken21 2 years, 2 months ago

Selected Answer: B

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/header-caching.html#header-caching-web-language>

upvoted 1 times

vherman 2 years, 3 months ago

Selected Answer: B

B is correct

upvoted 1 times

Steve_4542636 2 years, 3 months ago

Selected Answer: B

I think it's b

upvoted 2 times

LuckyAro 2 years, 3 months ago

Selected Answer: B

B is the correct answer

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #273

Topic 1


A rapidly growing ecommerce company is running its workloads in a single AWS Region. A solutions architect must create a disaster recovery (DR) strategy that includes a different AWS Region. The company wants its database to be up to date in the DR Region with the least possible latency. The remaining infrastructure in the DR Region needs to run at reduced capacity and must be able to scale up if necessary.

Which solution will meet these requirements with the LOWEST recovery time objective (RTO)?

- A. Use an Amazon Aurora global database with a pilot light deployment.
- B. Use an Amazon Aurora global database with a warm standby deployment. **Most Voted**
- C. Use an Amazon RDS Multi-AZ DB instance with a pilot light deployment.
- D. Use an Amazon RDS Multi-AZ DB instance with a warm standby deployment.

Correct Answer: B

Community vote distribution



Comments

Yechi **Highly Voted** 1 year, 9 months ago

Selected Answer: B

Note: The difference between pilot light and warm standby can sometimes be difficult to understand. Both include an environment in your DR Region with copies of your primary Region assets. The distinction is that pilot light cannot process requests without additional action taken first, whereas warm standby can handle traffic (at reduced capacity levels) immediately. The pilot light approach requires you to "turn on" servers, possibly deploy additional (non-core) infrastructure, and scale up, whereas warm standby only requires you to scale up (everything is already deployed and running). Use your RTO and RPO needs to help you choose between these approaches.

<https://docs.aws.amazon.com/whitepapers/latest/disaster-recovery-workloads-on-aws/disaster-recovery-options-in-the-cloud.html>

upvoted 28 times

nickolaj **Highly Voted** 1 year, 9 months ago

Selected Answer: B

Option A is incorrect because while Amazon Aurora global database is a good solution for disaster recovery, pilot light deployment provides only a minimalistic setup and would require manual intervention to make the DR Region fully

deployment provides only a minimalistic setup and would require manual intervention to make the DR Region fully operational, which increases the recovery time.

Option B is a better choice than Option A as it provides a warm standby deployment, which is an automated and more scalable setup than pilot light deployment. In this setup, the database is replicated to the DR Region, and the standby instance can be brought up quickly in case of a disaster.

Option C is incorrect because Multi-AZ DB instances provide high availability, not disaster recovery.

Option D is a good choice for high availability, but it does not meet the requirement for DR in a different region with the least possible latency.

upvoted 23 times

awsgeek75 Most Recent 10 months, 3 weeks ago

Selected Answer: B

B: Warm Standby is better when it comes to LOWEST RTO.

<https://docs.aws.amazon.com/whitepapers/latest/disaster-recovery-workloads-on-aws/disaster-recovery-options-in-the-cloud.html>

upvoted 2 times

pentium75 11 months, 2 weeks ago

Selected Answer: B

"Different Region" rules out C and D ("Multi-AZ" is within a region)

"Run at reduced capacity" = warm standby (while "pilot light" means that DR resources are shut down and are started manually in case of failover)

upvoted 4 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: B

The warm standby approach involves ensuring that there is a scaled down, but fully functional, copy of your production environment in another Region.

With the pilot light approach, you replicate your data from one Region to another and provision a copy of your core workload infrastructure. Resources required to support data replication and backup, such as databases and object storage, are always on. Other elements, such as application servers, are loaded with application code and configurations, but are "switched off".

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

An Amazon Aurora global database with a warm standby deployment provides continuous replication from one AWS Region to another, keeping the DR database up-to-date with minimal latency.

upvoted 2 times

A1975 1 year, 4 months ago

Selected Answer: B

In a Pilot Light scenario, only an EC2 Instance and a DB may be running. In Warm Standby, however, everything is running — in a much smaller capacity. This means the load balancer, gateways, databases, all subnets, and everything else are ready to go on a moment's notice.

with reference to below statement Option B is a better choice than Option A.

"The remaining infrastructure in the DR Region needs to run at reduced capacity and must be able to scale up if necessary".

upvoted 2 times

krisfromtw 1 year, 9 months ago

Selected Answer: D

should be D.

upvoted 2 times

pentium75 11 months, 2 weeks ago

"Multi-AZ" = multiple AZs in same region, but requirement is "a different AWS Region".

upvoted 3 times

leoattf 1 year, 9 months ago

No, my friend. The question asks for deployment in another Region. Hence, it cannot be C or D.
The answer is B because is Global (different regions) and Ward Standby has faster RTO than Pilot Light.
upvoted 10 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #274

Topic 1

A company runs an application on Amazon EC2 instances. The company needs to implement a disaster recovery (DR) solution for the application. The DR solution needs to have a recovery time objective (RTO) of less than 4 hours. The DR solution also needs to use the fewest possible AWS resources during normal operations.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. Create Amazon Machine Images (AMIs) to back up the EC2 instances. Copy the AMIs to a secondary AWS Region. Automate infrastructure deployment in the secondary Region by using AWS Lambda and custom scripts.
- B. Create Amazon Machine Images (AMIs) to back up the EC2 instances. Copy the AMIs to a secondary AWS Region. Automate infrastructure deployment in the secondary Region by using AWS CloudFormation. **Most Voted**
- C. Launch EC2 instances in a secondary AWS Region. Keep the EC2 instances in the secondary Region active at all times.
- D. Launch EC2 instances in a secondary Availability Zone. Keep the EC2 instances in the secondary Availability Zone active at all times.

Correct Answer: B*Community vote distribution*B (100%)**Comments****NolaH0la** **Highly Voted** 2 years, 3 months ago

Guys, sorry but I don't really have time to deepdive as my exam is soon. Based on chatGPT and my previous study the answer should be B

"Create Amazon Machine Images (AMIs) to back up the EC2 instances. Copy the AMIs to a secondary AWS Region. Automate infrastructure deployment in the secondary Region by using AWS CloudFormation," would likely be the most suitable solution for the given requirements.

This option allows for the creation of Amazon Machine Images (AMIs) to back up the EC2 instances, which can then be copied to a secondary AWS region to provide disaster recovery capabilities. The infrastructure deployment in the secondary region can be automated using AWS CloudFormation, which can help to reduce the amount of time and resources needed for deployment and management.

upvoted 11 times

NBone 1 year, 10 months ago

please how do you use chatGPT to study for these questions?

upvoted 5 times

nickolaj Highly Voted 2 years, 3 months ago

Selected Answer: B

Option B would be the most operationally efficient solution for implementing a DR solution for the application, meeting the requirement of an RTO of less than 4 hours and using the fewest possible AWS resources during normal operations.

By creating Amazon Machine Images (AMIs) to back up the EC2 instances and copying them to a secondary AWS Region, the company can ensure that they have a reliable backup in the event of a disaster. By using AWS CloudFormation to automate infrastructure deployment in the secondary Region, the company can minimize the amount of time and effort required to set up the DR solution.

upvoted 9 times

djgodzilla Most Recent 1 year, 5 months ago

Option E : Automate infrastructure deployment in the secondary Region by using terraform and ditch AWS CloudFormation ☐☐.

upvoted 6 times

pentium75 1 year, 5 months ago

Selected Answer: B

A is not "most operationally efficient"

C and D do not meet the "use the fewest possible AWS resources during normal operations" requirement

upvoted 3 times

vijaykamal 1 year, 8 months ago

Selected Answer: B

Option D suggests launching EC2 instances in a secondary Availability Zone (AZ), but AZs are not separate AWS Regions. While it provides high availability within a Region, it doesn't offer geographic redundancy, which is essential for disaster recovery.

upvoted 3 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: B

needs to use the fewest possible AWS resources during normal operations = backup & restore

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

Create Amazon Machine Images (AMIs) to back up the EC2 instances. Copy the AMIs to a secondary AWS Region. Automate infrastructure deployment in the secondary Region by using AWS CloudFormation

upvoted 2 times

AMYYMY 1 year, 9 months ago

B SHOULD BE RIGHT

upvoted 2 times

A1975 1 year, 10 months ago

Selected Answer: B

Option A: Add complexity and management overhead.

Option B: Creating AMIs for backup and using AWS CloudFormation for infrastructure deployment in the secondary Region is a more streamlined and automated approach. CloudFormation allows you to define and provision resources in a declarative manner, making it easier to maintain and update your infrastructure. This solution is more operationally efficient compared to Option A.

Option C: could be expensive and not fully aligned with the requirement of using the fewest possible AWS resources during normal operations.

Option D: might not be sufficient for meeting the DR requirements, as Availability Zones are still within the same AWS Region and might be subject to the same regional-level failures.

upvoted 2 times

NBone 1 year, 10 months ago

Please I would really appreciate clarification with this question. The community has voted 100% that the right answer is B. However, option D is shown to be the correct answer. So, who sets the correct answer? Which one should new comers like myself believe? the community's or the other (which am guessing is set by the moderators???) Please help.

upvoted 2 times

Kaula 1 year, 2 months ago

From Examtopics support:

In very few cases a conflict is raised between the provided and most-voted answers.

In case of conflict between the provided & most voted answers, we suggest our customers rely on the most voted answers and consider them most correct.

upvoted 2 times

SimiTik 2 years, 1 month ago

C may satisfy the requirement of using the fewest possible AWS resources during normal operations, it may not be the most operationally efficient or cost-effective solution in the long term.

upvoted 3 times

AlmeroSenior 2 years, 3 months ago

So Weird , they have product for this > Elastic Disaster Recovery , but option is not given .

upvoted 2 times

Yechi 2 years, 3 months ago

Selected Answer: B

https://docs.aws.amazon.com/zh_cn/whitepapers/latest/disaster-recovery-workloads-on-aws/disaster-recovery-options-in-the-cloud.html#backup-and-restore

upvoted 5 times

Joan111edu 2 years, 3 months ago

Selected Answer: B

the answer should be B

--->recovery time objective (RTO) of less than 4 hours.

https://docs.aws.amazon.com/zh_cn/whitepapers/latest/disaster-recovery-workloads-on-aws/disaster-recovery-options-in-the-cloud.html#backup-and-restore

upvoted 4 times

JA2018 6 months, 3 weeks ago

English version: <https://docs.aws.amazon.com/whitepapers/latest/disaster-recovery-workloads-on-aws/disaster-recovery-options-in-the-cloud.html>

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #275

Topic 1


A company runs an internal browser-based application. The application runs on Amazon EC2 instances behind an Application Load Balancer. The instances run in an Amazon EC2 Auto Scaling group across multiple Availability Zones. The Auto Scaling group scales up to 20 instances during work hours, but scales down to 2 instances overnight. Staff are complaining that the application is very slow when the day begins, although it runs well by mid-morning.

How should the scaling be changed to address the staff complaints and keep costs to a minimum?

- A. Implement a scheduled action that sets the desired capacity to 20 shortly before the office opens.
- B. Implement a step scaling action triggered at a lower CPU threshold, and decrease the cooldown period.
- C. Implement a target tracking action triggered at a lower CPU threshold, and decrease the cooldown period. **Most Voted**
- D. Implement a scheduled action that sets the minimum and maximum capacity to 20 shortly before the office opens.

Correct Answer: C

Community vote distribution



Comments

asoli **Highly Voted** 2 years, 2 months ago

Selected Answer: C

At first, I thought the answer is A. But it is C.

It seems that there is no information in the question about CPU or Memory usage.

So, we might think the answer is A. why? because what we need is to have the required (desired) number of instances. It already has scheduled scaling that works well in this scenario. Scale down after working hours and scale up in working hours. So, it just needs to adjust the desired number to start from 20 instances.

But here is the point it shows A is WRONG!!!

If it started with desired 20 instances, it will keep it for the whole day. What if the load is reduced? We do not need to keep the 20 instances always. That 20 is the MAXIMUM number we need, no the DESIRE number. So it is against COST that is the main objective of this question.

So, the answer is C

upvoted 33 times

c10356a 1 year, 5 months ago

There is no cooldown period in target tracking, but warm-up time.

upvoted 2 times

pentium75 1 year, 5 months ago

There is a cooldown period in the auto-scaling group, which helps 'keeping costs to a minimum' as instances would be removed sooner.

upvoted 4 times

mandragon 2 years, 1 month ago

If it starts with 20 instances it will not keep it all day. It will scale down based on demand. The scheduled action in Option A simply ensures that there are enough instances running to handle the increased traffic when the day begins, while still allowing the Auto Scaling group to scale up or down based on demand during the rest of the day.

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/scale-your-group.html>

upvoted 10 times

TheFivePips 1 year, 3 months ago

From what I can tell, you must specify an end time, or else it will run indefinitely. So I think A would be right, if they specified an end time. Otherwise C is more cost effective

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/ec2-auto-scaling-scheduled-scaling.html>

upvoted 3 times

xdkonorek2 1 year, 6 months ago

This is right, setting desired capacity doesn't turn off autoscaling policies

upvoted 3 times

meowruki Highly Voted 1 year, 6 months ago

Selected Answer: C

C. Implement a target tracking action triggered at a lower CPU threshold, and decrease the cooldown period.

Here's the reasoning:

Target Tracking Scaling Policy: With a target tracking scaling policy, you can set a target value for a specific metric, such as CPU utilization. The Auto Scaling group then adjusts the capacity to maintain that target.

Lower CPU Threshold: By triggering the target tracking action at a lower CPU threshold, the Auto Scaling group can proactively add instances when the workload increases, helping to address the slowness at the beginning of the day.

Decrease Cooldown Period: Reducing the cooldown period allows the Auto Scaling group to scale in and out more rapidly, making adjustments quicker in response to changing demand.

upvoted 6 times

meowruki 1 year, 6 months ago

Options A and D involve scheduled actions, which are time-based and may not be as responsive to immediate changes in demand. They also do not dynamically respond to varying workloads.

upvoted 2 times

srcntpc Most Recent 3 weeks, 5 days ago

Selected Answer: A

Why A is Correct:

Scheduled scaling action increases capacity to 20 instances shortly before office hours.

Ensures instances are already running and warm when employees start using the app.

This avoids delays caused by dynamic scaling and instance warm-up.

upvoted 2 times

zdi561 4 months ago

Selected Answer: A

Decreasing cpu threshold could cause more running instance(might have more than two at night) therefore higher cost. Schedule at desired to 20 minutes before the day will solve the issue that the app is very slow at beginning of the day

upvoted 2 times

upvoted 2 times

LeonSauveterre 6 months, 2 weeks ago

Selected Answer: A

Exam-wise, you gotta choose A.

The problem is explicitly described as slowness at the start of the day, which implies a predictable traffic pattern. Option C argues for a responsive approach (reacting dynamically to real-time demand changes) rather than a predictive approach (pre-scaling based on a known schedule).

Also, option A can minimize costs compared to option D as well, because the desired capacity is a flexible number that the Auto Scaling group targets at specific times but doesn't enforce as a hard minimum.

If the emphasis is on handling varying traffic patterns throughout the day, Option C could be considered viable.

upvoted 1 times

AbhiBK 9 months, 1 week ago

Question is about cost effectiveness hence Use target tracking scaling policies to maintain a specific metric, such as CPU utilization or request count per target. This allows the Auto Scaling group to dynamically adjust the number of instances based on real-time demand. We do not need to have 20 instances up and running during start of the day.

upvoted 2 times

maryam_sh 9 months, 2 weeks ago

Selected Answer: A

Not C because: target tracking scaling attempts to maintain a target metric (like average CPU utilization). While more responsive, it would still react to increased load rather than pre-scaling. The initial slow period would persist as the scaling reacts to the increased demand rather than anticipating it.

upvoted 2 times

ChymKuBoy 11 months, 1 week ago

Selected Answer: A

A for sure.

upvoted 3 times

foha2012 1 year, 4 months ago

Answer is A. Makes more sense to me.

upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: C

A is not cost effective, it would set the number of instances to maximum even before the first employee arrives.

D is not cost effective, it would cause the permanent use of 20 instances

B could almost work, but if you configure small steps then it scales too slowly in the morning; if you configure big steps (like "add 8 instances at a time") it would scale in the morning but not be cost-efficient during the day.

C would address the requirement, it would scale to meet a certain CPU utilization. Decreasing the cooldown period (which is not possible for the scaling policy itself but for the auto-scaling group) would help 'keeping costs to a minimum'.

upvoted 4 times

Mikado211 1 year, 5 months ago

Selected Answer: A

The question 369 is exactly the same problem,

Since a scheduled scaling doesn't disable the autoscaling later in the day the A works perfectly well.

upvoted 3 times

Cyberkayu 1 year, 5 months ago

A. since only a boot storm issue at 9am and settle down in mid morning, 20 instance is enough to support the workload NOT C. Reduce threshold to trigger (lets say 50% from 80% utilization) and lower cool down period, will still take time to ramp up to max 20 instance.

upvoted 1 times

pentium75 1 year, 5 months ago

How would scaling up to the maximum number of instances "keep costs to a minimum"?

upvoted 2 times

[Removed] 1 year, 6 months ago

Selected Answer: A

I would go with A. Autoscaling is still there and the problem is clearly in morning.

upvoted 4 times

pentium75 1 year, 5 months ago

WOuld not keep costs to a minimum

upvoted 2 times

wearrexdzw3123 1 year, 6 months ago

My mistake, I should have chosen c. A lower threshold can expand in advance, and lowering cooling can increase the expansion frequency.

upvoted 3 times

wearrexdzw3123 1 year, 6 months ago

Selected Answer: A

I choose option A because the root of the problem is the inability of the scaling speed in the morning to meet the demand, rather than what criteria to use for scaling.

upvoted 1 times

pentium75 1 year, 5 months ago

How would scaling up to the maximum number of instances "keep costs to a minimum"?

upvoted 2 times

TariqKipkemei 1 year, 8 months ago

To keep costs to a minimum target tracking is the best option.

For example the scaling metric is the average CPU utilization of the EC2 auto scaling instances, and their average during the day should always be 80%. When CloudWatch detects that the average CPU utilization is beyond 80% at start of day, it will trigger the target tracking policy to scale out the auto scaling group to meet this target utilization. Once everything is settled and the average CPU utilization has gone below 80% at night, another scale in action will kick in and reduce the number of auto scaling instances in the auto scaling group.

upvoted 4 times

TariqKipkemei 1 year, 8 months ago

Option C is best

upvoted 2 times

Ramdi1 1 year, 8 months ago

Selected Answer: A

I am going A based on it stating upto 20 so you already know what they maximum they use which is n a sense consistent. however i can see why people have put C. I think they need more clarification on the questions.

upvoted 3 times

pentium75 1 year, 5 months ago

How would scaling up to the maximum number of instances "keep costs to a minimum"?

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #276

Topic 1


A company has a multi-tier application deployed on several Amazon EC2 instances in an Auto Scaling group. An Amazon RDS for Oracle instance is the application's data layer that uses Oracle-specific PL/SQL functions. Traffic to the application has been steadily increasing. This is causing the EC2 instances to become overloaded and the RDS instance to run out of storage. The Auto Scaling group does not have any scaling metrics and defines the minimum healthy instance count only. The company predicts that traffic will continue to increase at a steady but unpredictable rate before leveling off.

What should a solutions architect do to ensure the system can automatically scale for the increased traffic? (Choose two.)

- A. Configure storage Auto Scaling on the RDS for Oracle instance. **Most Voted**
- B. Migrate the database to Amazon Aurora to use Auto Scaling storage.
- C. Configure an alarm on the RDS for Oracle instance for low free storage space.
- D. Configure the Auto Scaling group to use the average CPU as the scaling metric. **Most Voted**
- E. Configure the Auto Scaling group to use the average free memory as the scaling metric.

Correct Answer: AD

Community vote distribution



Comments

klayytech **Highly Voted** 2 years, 2 months ago

Selected Answer: AD

A) Configure storage Auto Scaling on the RDS for Oracle instance.
= Makes sense. With RDS Storage Auto Scaling, you simply set your desired maximum storage limit, and Auto Scaling takes care of the rest.

B) Migrate the database to Amazon Aurora to use Auto Scaling storage.
= Scenario specifies application's data layer uses Oracle-specific PL/SQL functions. This rules out migration to Aurora.

C) Configure an alarm on the RDS for Oracle instance for low free storage space.
= You could do this but what does it fix? Nothing. The CW notification isn't going to trigger anything.

D) Configure the Auto Scaling group to use the average CPU as the scaling metric.
= Makes sense. The CPU utilization is the precursor to the storage outage. When the ec2 instances are overloaded, the RDS instance storage hits its limits. too.

upvoted 29 times

[Removed] 1 year, 3 months ago

Nicely Explained.

upvoted 3 times

BombArat Most Recent 11 months, 2 weeks ago

Selected Answer: AE

B mean db migration and only oracle specific commands are allowed

C is just notification not High availability

Now the toss up is between D and E since D is measuring CPU % and E is measuring memory and the question states "This is causing the EC2 instances to become overloaded and the RDS instance to run out of storage" i will err on the side of option E

upvoted 1 times

Uzbekistan 1 year, 2 months ago

Selected Answer: AC

To ensure the system can automatically scale for the increased traffic, you can take the following steps:

A. Configure storage Auto Scaling on the RDS for Oracle instance.

By enabling storage Auto Scaling on the RDS instance, you ensure that additional storage is provisioned automatically when the existing storage reaches capacity. This helps prevent the RDS instance from running out of storage due to increased traffic and data growth.

C. Configure an alarm on the RDS for Oracle instance for low free storage space.

Setting up an alarm for low free storage space on the RDS instance allows you to receive notifications when the storage capacity is approaching its limits. This proactive monitoring helps you take necessary actions, such as adding more storage or scaling resources, before it affects the application's performance.

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: AD

B is not possible because the application "uses Oracle-specific PL/SQL functions"

C does not meet the "automatically scale" requirement

E would require an agent on the hosts which we might not have, plus CPU is a better indicator than memory

upvoted 2 times

meowruki 1 year, 6 months ago

Selected Answer: AD

A. Configure storage Auto Scaling on the RDS for Oracle instance.

This option allows the RDS instance to automatically scale its storage based on the actual storage usage, ensuring that you don't run out of storage.

D. Configure the Auto Scaling group to use the average CPU as the scaling metric.

By using CPU utilization as a scaling metric, the Auto Scaling group can dynamically adjust the number of EC2 instances based on the application's demand. This helps in handling increased traffic and preventing overload on existing instances.

upvoted 2 times

meowruki 1 year, 6 months ago

Option B (Migrate the database to Amazon Aurora): While Amazon Aurora provides benefits such as auto-scaling storage and high performance, it involves migrating from Oracle to Aurora, which might require application changes and data migration efforts.

Option C (Configure an alarm on the RDS for Oracle instance for low free storage space): While it's good to have an alarm for low storage space, configuring storage Auto Scaling (Option A) is a more proactive solution that automatically adjusts storage before reaching a critical point.

Option E (Configure the Auto Scaling group to use the average free memory as the scaling metric): While monitoring memory is important for application performance, CPU utilization is often a more direct and responsive metric for auto-scaling in many scenarios.

upvoted 3 times

JA2018 6 months, 3 weeks ago

STEM mentioned the application "uses Oracle-specific PL/SQL functions"
upvoted 1 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: AD

Configure storage Auto Scaling on the RDS for Oracle instance and Configure the Auto Scaling group to use the average CPU as the scaling metric to accommodate the increased traffic automatically.

upvoted 2 times

vijaykamal 1 year, 8 months ago

Selected Answer: AD

Option B (Migrate the database to Amazon Aurora) may be a good long-term solution, but it involves database migration, which can be complex and time-consuming. For immediate scalability and to address the storage issue, configuring storage Auto Scaling on the existing RDS instance is a more immediate and straightforward solution.

Option C (Configure an alarm on the RDS for Oracle instance for low free storage space) is useful for monitoring, but it doesn't proactively address the storage issue by automatically expanding storage as needed.

Option E (Configure the Auto Scaling group to use the average free memory as the scaling metric) is less common as a scaling metric for EC2 instances compared to CPU utilization. While memory can be an important factor for application performance, CPU utilization is typically a more commonly used metric for scaling decisions. It also doesn't directly address the RDS storage issue.

upvoted 2 times

JA2018 6 months, 3 weeks ago

the application "uses Oracle-specific PL/SQL functions"

Rules out Option B in my mind
upvoted 1 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: AD

A. By enabling storage Auto Scaling on the RDS for Oracle instance, it will automatically add more storage when the existing storage is running out, ensuring the application's data layer can handle the increased data storage requirements.

D. By configuring the Auto Scaling group to use the average CPU utilization as the scaling metric, it can automatically add more EC2 instances to the Auto Scaling group when the CPU utilization exceeds a certain threshold. This will help handle the increased traffic and workload on the EC2 instances in the multi-tier application.

upvoted 2 times

A1975 1 year, 10 months ago

Selected Answer: AD

A. By enabling storage Auto Scaling on the RDS for Oracle instance, it will automatically add more storage when the existing storage is running out, ensuring the application's data layer can handle the increased data storage requirements.

D. By configuring the Auto Scaling group to use the average CPU utilization as the scaling metric, it can automatically add more EC2 instances to the Auto Scaling group when the CPU utilization exceeds a certain threshold. This will help handle the increased traffic and workload on the EC2 instances in the multi-tier application.

upvoted 2 times

kruasan 2 years, 1 month ago

Selected Answer: AD

These options will allow the system to scale both the compute tier (EC2 instances) and the data tier (RDS storage) automatically as traffic increases:

A. Storage Auto Scaling will allow the RDS for Oracle instance to automatically increase its allocated storage when free storage space gets low. This ensures the database does not run out of capacity and can continue serving data to the application.

D. Configuring the EC2 Auto Scaling group to scale based on average CPU utilization will allow it to launch additional instances automatically as traffic causes higher CPU levels across the instances. This scales the compute tier to handle increased demand.

upvoted 3 times

kraken21 2 years, 2 months ago

Selected Answer: AD

Auto scaling storage RDS will ease storage issues and migrating Oracle PL/Sql to Aurora is cumbersome. Also Aurora has auto storage scaling by default.

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_PIOPS.StorageTypes.html#USER_PIOPS.Autoscaling

upvoted 3 times

Nel8 2 years, 3 months ago

Selected Answer: BD

My answer is B & D...

B. Migrate the database to Amazon Aurora to use Auto Scaling Storage. --- Aurora storage is also self-healing. Data blocks and disks are continuously scanned for errors and repaired automatically.

D. Configure the Auto Scaling group to sue the average CPU as the scaling metric. -- Good choice.

I believe either A & C or B & D options will work.

upvoted 3 times

FourOfAKind 2 years, 3 months ago

In this question, you have Oracle DB, and Amazon Aurora is for MySQL/PostgreSQL. A and D are the correct choices.

upvoted 5 times

[Removed] 2 years, 2 months ago

You can migrate Oracle PL/SQL to Aurora:

<https://docs.aws.amazon.com/dms/latest/oracle-to-aurora-mysql-migration-playbook/chap-oracle-aurora-mysql.sql.html>

upvoted 1 times

pentium75 1 year, 5 months ago

But the application "uses Oracle-specific PL/SQL functions".

upvoted 2 times

[Removed] 2 years, 2 months ago

I still think A is the answer, because RDS for Oracle auto scaling once enabled it will automatically adjust the storage capacity.

upvoted 2 times

Ja13 2 years, 3 months ago

Selected Answer: AD

a and d

upvoted 4 times

KZM 2 years, 3 months ago

A and D.

upvoted 4 times

GwonLEE 2 years, 3 months ago

Selected Answer: AD

a and d

upvoted 4 times

LuckyAro 2 years, 3 months ago

Selected Answer: AD

A and D

upvoted 3 times

Joan111edu 2 years, 3 months ago

Selected Answer: AD

<https://www.examtopics.com/discussions/amazon/view/46534-exam-aws-certified-solutions-architect-associate-saa-c02/>
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #277

Topic 1

A company provides an online service for posting video content and transcoding it for use by any mobile platform. The application architecture uses Amazon Elastic File System (Amazon EFS) Standard to collect and store the videos so that multiple Amazon EC2 Linux instances can access the video content for processing. As the popularity of the service has grown over time, the storage costs have become too expensive.

Which storage solution is MOST cost-effective?

- A. Use AWS Storage Gateway for files to store and process the video content.
- B. Use AWS Storage Gateway for volumes to store and process the video content.
- C. Use Amazon EFS for storing the video content. Once processing is complete, transfer the files to Amazon Elastic Block Store (Amazon EBS).
- D. Use Amazon S3 for storing the video content. Move the files temporarily over to an Amazon Elastic Block Store (Amazon EBS) volume attached to the server for processing. **Most Voted**

Correct Answer: D

Community vote distribution

D (82%)

A (18%)

Comments

bdp123 **Highly Voted** 2 years, 3 months ago

Selected Answer: D

Storage gateway is not used for storing content - only to transfer to the Cloud
upvoted 28 times

pentium75 1 year, 5 months ago

A doesn't say "store content ON the gateway", it says "use AWS Storage Gateway for files" (which is the product) "to store the video content" [on S3].

upvoted 3 times

jaswantn 1 year, 4 months ago

Creating storage gateway for Files will mount S3 bucket as an NFS volume that can be shared among EC2 Instances in the

same manner as EFS but more cost effectively.

upvoted 1 times

kraken21 Highly Voted 2 years, 2 months ago

Selected Answer: D

There is no on-prem/non Aws infrastructure to create a gateway. Also, EFS+EBS is more expensive than EFS and S3. So D is the best option.

upvoted 12 times

pentium75 1 year, 5 months ago

But how do you attach "an EBS volume" to all the servers, and how will you use the files on it then to serve customers.

upvoted 3 times

FlyingHawk Most Recent 6 months, 1 week ago

Selected Answer: D

Option D (S3 + EBS) is the most cost-effective and scalable choice for a service that involves storing video content and performing I/O-intensive transcoding for mobile platforms.

Option A (S3 File Gateway) could work for applications with strict file system requirements but is less optimal for high-performance transcoding which is required by this online service

upvoted 2 times

Uzbekistan 1 year, 2 months ago

Selected Answer: D

D. Use Amazon S3 for storing the video content. Move the files temporarily over to an Amazon Elastic Block Store (Amazon EBS) volume attached to the server for processing.

Amazon S3 (Simple Storage Service) is highly durable, scalable, and cost-effective for storing large volumes of data, such as video content. It offers lower storage costs compared to Amazon EFS and is suitable for storing large files like video content.

For processing the video content, you can temporarily move the files from Amazon S3 to an Amazon EBS volume attached to the EC2 instances. This approach allows you to leverage the high-performance storage of Amazon EBS for processing, while still benefiting from the cost-effectiveness of Amazon S3 for long-term storage. Once processing is complete, you can remove the temporary files from the EBS volume and store the final results back in S3.

upvoted 4 times

bujuman 1 year, 3 months ago

Selected Answer: A

Answer is closer to the following principle and D is near impossible to implement:

"Amazon S3 File Gateway – Amazon S3 File Gateway supports a file interface into Amazon Simple Storage Service (Amazon S3) and combines a service and a virtual software appliance. By using this combination, you can store and retrieve objects in Amazon S3 using industry-standard file protocols such as Network File System (NFS) and Server Message Block (SMB). You deploy the gateway into your on-premises environment as a virtual machine (VM) running on VMware ESXi, Microsoft Hyper-V, or Linux Kernel-based Virtual Machine (KVM), or as a hardware appliance that you order from your preferred reseller. You can also deploy the Storage Gateway VM in VMware Cloud on AWS, or as an AML in Amazon EC2. The gateway provides access to objects in S3 as files or file share mount points. With a S3 File Gateway, you can do the following"

upvoted 1 times

vip2 1 year, 3 months ago

Selected Answer: A

A is correct

For D, how to move file from S3 to EBS temporarily????

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

I was initially going for D but EBS part makes no sense as it is not possible. Closest explanation of A is in this article:

<https://aws.amazon.com/blogs/storage/mounting-amazon-s3-to-an-amazon-ec2-instance-using-a-private-connection-to-s3-file-gateway/>

A is missing a lot of key steps but D is just impossible. Maybe it's just the wording?

upvoted 1 times

AzExam2020 1 year, 5 months ago

EFS is already used, why EBS is an option in the answer?

upvoted 2 times

anikolov 1 year, 5 months ago

Selected Answer: A

AWS Storage Gateway S3 file gateway can be setup on EC2 (<https://repost.aws/knowledge-center/file-gateway-ec2>). It use local disks/EBS for caching data.

D: Can be used too, using attached EBS volume to each EC2 instance to process files. If require single EBS volume to be attached to multiple EC2, then it is possible too if they are in the same Availability Zone ->

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html>.

For me A and D both are possible, but expect AWS would like to select Storage GW

upvoted 2 times

SVDK 1 year, 4 months ago

I agree. This documentation has convinced me. <https://docs.aws.amazon.com/filegateway/latest/files3/what-is-file-s3.html>

upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: A

A would work, "Storage Gateway for files" can provide access to S3 (cheap) via NFS (what the clients are using now). It has some additional cost in addition to the S3 charges, but would still be way cheaper than EFS.

B would work for a single server, but as it provides a volume via iSCSI, it could be mounted only to a single server - does not meet the 'multiple instances can access' requirement.

C and D do not meet the 'multiple instances can access' requirement because EBS can't be easily attached to all servers at the same time.

upvoted 2 times

pentium75 1 year, 5 months ago

And even if ignoring the 'multiple instances can access' requirement, D would be against WAF; for temporary storage you'd use instance storage, not EBS.

upvoted 2 times

liux99 1 year, 7 months ago

Storage gateway is intended for on-premises applications to access cloud storage, so A, B is out. The question explicitly states that the files are uploaded and stored in EFS, not S3, so D is not correct. The answer is C. The EFS storage costs 10 times more than EBS, so moving files to EBS after processing is the solution.

upvoted 1 times

beginnercloud 1 year, 7 months ago

Selected Answer: D

Answer D is correct.

Storage gateway is not used for storing content - only to transfer to the Cloud

upvoted 2 times

pentium75 1 year, 5 months ago

But how do you attach "an EBS volume" to all the servers, and how will you use the files on it then to serve customers?

A doesn't say "store content ON the gateway", it says "use AWS Storage Gateway for files" (which is the product) "to store the video content" [on S3]. And to be exact, Storage Gateway is not "to transfer to the cloud" but to provide access to S3 storage via SMB or NFS.

upvoted 1 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: D

Cost effective = Use Amazon S3 for storing the video content. Move the files temporarily over to an Amazon Elastic Block Store (Amazon EBS) volume attached to the server for processing

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

Amazon S3 provides low-cost object storage for storing large amounts of unstructured data like videos. The videos can be stored in S3 durably and reliably.

For processing, the video files can be temporarily copied from S3 to an EBS volume attached to the EC2 instance. EBS provides low latency block storage for high performance video processing.

Once processing is complete, the output can be stored back in S3.

upvoted 5 times

bjexamprep 1 year, 10 months ago

Selected Answer: D

The question doesn't give enough information. Well, quite a few AWS exam questions don't provide enough info. Ideally, A could be the best answer if it mentions S3 as the backend of storage gateway. Because if it doesn't mention S3 as the backend, that implies either Storage gateway as the storage(which is impossible) or continue using EFS(also impossible). D is not ideal, because it will introduce video download cost for downloading files from S3 to EBS temporary storage. But it is the best option we have.

upvoted 2 times

pentium75 1 year, 5 months ago

A mentions "AWS Storage Gateway for files" which implies S3 as the backend storage. D does not meet the 'multiple instances can access for processing' requirement.

upvoted 1 times

foha2012 1 year, 4 months ago

We are ditching EFS in favor of S3. So there is no longer simultaneous access happening. Whoever needs the file, downloads it from S3, process it on their EC2 instance and save it back to S3.

upvoted 1 times

Undisputed 1 year, 10 months ago

Selected Answer: D

A more cost-effective storage solution for this scenario would be Amazon Simple Storage Service (Amazon S3). Amazon S3 is an object storage service that offers high scalability, durability, and availability at a lower cost compared to Amazon EFS. By using Amazon S3, you only pay for the storage you use, and it is typically more cost-efficient for scenarios where data is accessed less frequently, such as video storage for processing.

upvoted 3 times

smartegnine 1 year, 11 months ago

Selected Answer: A

The result should be A.

Amazon storage gateway has 4 types, S3 File Gateway, FSx file gateway, Type Gateway and Volume Gateway.

If not specific reference file gateway should be default as S3 gateway, which sent file over to S3 the most cost effective storage in AWS.

Why not D, the reason is last sentence, there are multiple EC2 servers for processing the video and EBS can only attach to 1 EC2 instance at a time, so if you use EBS, which mean for each EC2 instance you will have 1 EBS. This rule out D.

upvoted 1 times

argl1995 1 year, 11 months ago

We can use multi-attach feature of EBS to attach one EBS volume to multiple Ec2 instances

upvoted 2 times

[Removed] 1 year, 11 months ago

AWS Storage Gateway = extend storage to onprem

upvoted 1 times

pentium75 1 year, 5 months ago

Storage Gateway is a GATEWAY, it does not store anything. You could use the gateway as a cache for content actually in S3. Btw, A and B even say that Storage Gateway would "process the video content" ...

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #278

Topic 1

A company wants to create an application to store employee data in a hierarchical structured relationship. The company needs a minimum-latency response to high-traffic queries for the employee data and must protect any sensitive data. The company also needs to receive monthly email messages if any financial information is present in the employee data.

Which combination of steps should a solutions architect take to meet these requirements? (Choose two.)

- A. Use Amazon Redshift to store the employee data in hierarchies. Unload the data to Amazon S3 every month.
- B. Use Amazon DynamoDB to store the employee data in hierarchies. Export the data to Amazon S3 every month. **Most Voted**
- C. Configure Amazon Macie for the AWS account. Integrate Macie with Amazon EventBridge to send monthly events to AWS Lambda.
- D. Use Amazon Athena to analyze the employee data in Amazon S3. Integrate Athena with Amazon QuickSight to publish analysis dashboards and share the dashboards with users.
- E. Configure Amazon Macie for the AWS account. Integrate Macie with Amazon EventBridge to send monthly notifications through an Amazon Simple Notification Service (Amazon SNS) subscription. **Most Voted**

Correct Answer: BE

Community vote distribution

BE (100%)

Comments

Bhawesh **Highly Voted** 1 year, 9 months ago

Selected Answer: BE

Data in hierarchies : Amazon DynamoDB

B. Use Amazon DynamoDB to store the employee data in hierarchies. Export the data to Amazon S3 every month.

Sensitive Info: Amazon Macie

E. Configure Amazon Macie for the AWS account. Integrate Macie with Amazon EventBridge to send monthly notifications through an Amazon Simple Notification Service (Amazon SNS) subscription.

upvoted 17 times

gold4otas 1 year, 8 months ago

Can someone please provide explanation whv options "B" & "C" are the correct options?

Can someone please provide explanation why options B and E are the correct options.
upvoted 1 times

smartegnine 1 year, 5 months ago

C is half statement once event sent to Lambda what is next? Should send email right, but it does not say it.
upvoted 2 times

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: BE

B and E are the steps to meet all of the requirements.

B meets the need to store hierarchical employee data in DynamoDB for low latency queries at high traffic. DynamoDB can handle the access patterns for hierarchical data. Exporting to S3 monthly provides an audit trail.

E sets up Macie to analyze sensitive data and integrate with EventBridge to trigger monthly SNS notifications when financial data is present.

upvoted 6 times

hro **Most Recent** 8 months, 3 weeks ago

B - because to store employee data in a hierarchical structured relationship. AmazonDB "...Schema flexibility lets DynamoDB store complex hierarchical data within a single item."

E - because C omits the monthly email notifications resolved by using Amazon SNS.

Just my take.

upvoted 2 times

[Removed] 9 months, 2 weeks ago

Why not C and D? Can anyone explain please.

upvoted 1 times

Cyberkayu 12 months ago

A. Unload the data to Amazon S3 every month.

doesn't make sense to empty the employee data from redshift monthly

upvoted 2 times

beginnercloud 1 year, 1 month ago

Selected Answer: BE

B and E are the steps to meet all of the requirements.

upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: BE

Use Amazon DynamoDB to store the employee data in hierarchies. Export the data to Amazon S3 every month. Configure Amazon Macie for the AWS account. Integrate Macie with Amazon EventBridge to send monthly notifications through an Amazon Simple Notification Service (Amazon SNS) subscription.

upvoted 2 times

A1975 1 year, 4 months ago

Selected Answer: BE

JB. Amazon DynamoDB is a fully managed NoSQL database service that provides low-latency, high-performance storage for hierarchical data. It handles high-traffic queries and delivering fast responses to retrieve employee data efficiently.

E. Amazon Macie is a service that uses machine learning to automatically discover, classify, and protect sensitive data in AWS. Integrating Macie with Amazon EventBridge allows you to receive events whenever any financial information is identified in the employee data. By using Amazon SNS, you can receive these notifications via email.

upvoted 3 times

cesargalindo123 1 year, 5 months ago

AE

<https://aws.amazon.com/es/blogs/big-data/query-hierarchical-data-models-within-amazon-redshift/>

upvoted 3 times

kruasan 1 year, 7 months ago

Selected Answer: BE

, the combination of DynamoDB for fast data queries, S3 for durable storage and backups, Macie for sensitive data monitoring, and EventBridge + SNS for email notifications satisfies all needs: fast query response, sensitive data protection, and monthly alerts. The solutions architect should implement DynamoDB with export to S3, and configure Macie with integration to send SNS email notifications.

upvoted 2 times

kruasan 1 year, 7 months ago

Generally, for building a hierarchical relationship model, a graph database such as Amazon Neptune is a better choice. In some cases, however, DynamoDB is a better choice for hierarchical data modeling because of its flexibility, security, performance, and scale.

<https://docs.aws.amazon.com/prescriptive-guidance/latest/dynamodb-hierarchical-data-model/introduction.html>

upvoted 4 times

darn 1 year, 7 months ago

why Dynamo and not Redshift?

upvoted 2 times

kruasan 1 year, 7 months ago

1. Low latency - DynamoDB is a fully managed NoSQL database service that provides fast and predictable performance with single-digit millisecond latency. Redshift is a data warehouse solution optimized for complex analytical queries, so query latency would typically be higher. Since the requirements specify minimum latency for high-traffic queries, DynamoDB is better suited.

2. Scalability - DynamoDB is highly scalable, able to handle very high read and write throughput with no downtime. Redshift also scales, but may experience some downtime during rescale operations. For a high-traffic application, DynamoDB's scalability and availability are better matched.

upvoted 5 times

kruasan 1 year, 7 months ago

3. Hierarchical data - DynamoDB supports hierarchical (nested) data structures well in a NoSQL data model. Defining hierarchical employee data may be more complex in Redshift's columnar SQL data warehouse structure. DynamoDB is built around flexible data schemas that can represent complex relationships.

4. Data export - Both DynamoDB and Redshift allow exporting data to S3, so that requirement could be met with either service. However, overall DynamoDB is the better fit based on the points above regarding latency, scalability, and support for hierarchical data.

upvoted 6 times

PRASAD180 1 year, 9 months ago

BE is crt 100%

upvoted 2 times

KZM 1 year, 9 months ago

B and E

To send monthly email messages, an SNS service is required.

upvoted 3 times

skiwili 1 year, 9 months ago

Selected Answer: BE

B and E

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #279

Topic 1

A company has an application that is backed by an Amazon DynamoDB table. The company's compliance requirements specify that database backups must be taken every month, must be available for 6 months, and must be retained for 7 years.

Which solution will meet these requirements?

- A. Create an AWS Backup plan to back up the DynamoDB table on the first day of each month. Specify a lifecycle policy that transitions the backup to cold storage after 6 months. Set the retention period for each backup to 7 years. **Most Voted**
- B. Create a DynamoDB on-demand backup of the DynamoDB table on the first day of each month. Transition the backup to Amazon S3 Glacier Flexible Retrieval after 6 months. Create an S3 Lifecycle policy to delete backups that are older than 7 years.
- C. Use the AWS SDK to develop a script that creates an on-demand backup of the DynamoDB table. Set up an Amazon EventBridge rule that runs the script on the first day of each month. Create a second script that will run on the second day of each month to transition DynamoDB backups that are older than 6 months to cold storage and to delete backups that are older than 7 years.
- D. Use the AWS CLI to create an on-demand backup of the DynamoDB table. Set up an Amazon EventBridge rule that runs the command on the first day of each month with a cron expression. Specify in the command to transition the backups to cold storage after 6 months and to delete the backups after 7 years.

Correct Answer: A*Community vote distribution***Comments****vijaykamal** **Highly Voted** 1 year, 8 months ago**Selected Answer: A**

Option B mentions using Amazon S3 Glacier Flexible Retrieval, but DynamoDB doesn't natively support transitioning backups to Amazon S3 Glacier. Options C and D involve custom scripts and EventBridge rules, which add complexity and may not be as reliable or efficient as using AWS Backup for this purpose.

upvoted 14 times

MNotABot **Highly Voted** 1 year, 11 months ago

All except A are "On-demand"

upvoted 6 times

ChymKuBoy Most Recent 11 months, 1 week ago

Selected Answer: A

A for sure

upvoted 2 times

Sadiya_Javid_Abbasi 1 year, 5 months ago

Selected Answer: A

<https://docs.aws.amazon.com/aws-backup/latest/devguide/creating-a-backup-plan.html>

upvoted 2 times

mwwt2022 1 year, 5 months ago

Why B is wrong?

upvoted 1 times

Cyberkayu 1 year, 5 months ago

BCD, on-demand backup, manual work

upvoted 5 times

beginnercloud 1 year, 7 months ago

Selected Answer: A

<https://aws.amazon.com/blogs/database/set-up-scheduled-backups-for-amazon-dynamodb-using-aws-backup/>

upvoted 4 times

chanchal133 1 year, 9 months ago

Selected Answer: A

A is right ans

upvoted 2 times

narddrer 1 year, 11 months ago

Selected Answer: B

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/BackupRestore.html>

Using DynamoDB with AWS Backup, you can copy your on-demand backups across AWS accounts and Regions, add cost allocation tags to on-demand backups, and transition on-demand backups to cold storage for lower costs. To use these advanced features, you must opt in to AWS Backup.

upvoted 4 times

pentium75 1 year, 5 months ago

"on demand" (manual) backup -> hell no

upvoted 4 times

kruasan 2 years, 1 month ago

Selected Answer: A

This solution satisfies the requirements in the following ways:

- AWS Backup will automatically take full backups of the DynamoDB table on the schedule defined in the backup plan (the first of each month).
- The lifecycle policy can transition backups to cold storage after 6 months, meeting that requirement.
- Setting a 7-year retention period in the backup plan will ensure each backup is retained for 7 years as required.
- AWS Backup manages the backup jobs and lifecycle policies, requiring no custom scripting or management.

upvoted 3 times

TariqKipkemei 2 years, 2 months ago

Answer is A

upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Create an AWS Backup plan to back up the DynamoDB table on the first day of each month. Specify a lifecycle policy that transitions the backup to cold storage after 6 months. Set the retention period for each backup to 7 years

upvoted 2 times

mmustafa4455 2 years, 2 months ago

Selected Answer: A

The correct Answer is A

<https://aws.amazon.com/blogs/database/set-up-scheduled-backups-for-amazon-dynamodb-using-aws-backup/>

upvoted 2 times

mmustafa4455 2 years, 2 months ago

Its B.

<https://aws.amazon.com/blogs/database/set-up-scheduled-backups-for-amazon-dynamodb-using-aws-backup/>

upvoted 2 times

Wael216 2 years, 3 months ago

Selected Answer: A

A is the answer

upvoted 2 times

LuckyAro 2 years, 3 months ago

Selected Answer: A

A is the answer.

upvoted 2 times

skiwili 2 years, 3 months ago

Selected Answer: A

A is the correct answer

upvoted 1 times

rrharris 2 years, 3 months ago

A is the Answer

can be used to create backup schedules and retention policies for DynamoDB tables

upvoted 3 times

kpato87 2 years, 3 months ago

Selected Answer: A

A. Create an AWS Backup plan to back up the DynamoDB table on the first day of each month. Specify a lifecycle policy that transitions the backup to cold storage after 6 months. Set the retention period for each backup to 7 years.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #280

Topic 1

A company is using Amazon CloudFront with its website. The company has enabled logging on the CloudFront distribution, and logs are saved in one of the company's Amazon S3 buckets. The company needs to perform advanced analyses on the logs and build visualizations.

What should a solutions architect do to meet these requirements?

- A. Use standard SQL queries in Amazon Athena to analyze the CloudFront logs in the S3 bucket. Visualize the results with AWS Glue.
- B. Use standard SQL queries in Amazon Athena to analyze the CloudFront logs in the S3 bucket. Visualize the results with Amazon QuickSight. **Most Voted**
- C. Use standard SQL queries in Amazon DynamoDB to analyze the CloudFront logs in the S3 bucket. Visualize the results with AWS Glue.
- D. Use standard SQL queries in Amazon DynamoDB to analyze the CloudFront logs in the S3 bucket. Visualize the results with Amazon QuickSight.

Correct Answer: B*Community vote distribution*

B (95%)

D (5%)

Comments**pentium75** **Highly Voted** 11 months, 2 weeks ago**Selected Answer: B**

Data is in S3 -> Athena, not DynamoDB (thus A or B)
Visualize -> QuickSight, not Glue (thus B or D)
upvoted 9 times

rrharris **Highly Voted** 1 year, 9 months ago

Answer is B - Quicksite creating data visualizations

<https://docs.aws.amazon.com/quicksight/latest/user/welcome.html>
upvoted 8 times

LoXoL **Most Recent** 10 months, 2 weeks ago

Selected Answer: B

Glue is meant to prepare and transform data for analytics, not to build visualizations. Hence A and C are out. Athena is used to analyze data stored in S3 and it is commonly used with QuickSight, thus B is the answer
upvoted 4 times

[Removed] 12 months ago

Admin please remove my comment. That answer was for another question.
upvoted 1 times

[Removed] 12 months ago

A & C combined make sense isn't it?
upvoted 1 times

pentium75 11 months, 2 weeks ago

Why 'use queries in DynamoDB" when the data is in S3? And why Glue?
upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

OptionB: Amazon Athena allows you to run standard SQL queries directly on the data stored in the S3 bucket. Amazon QuickSight is a business intelligence (BI) service that allows you to create interactive and visual dashboards to analyze data. You can connect Amazon QuickSight to Amazon Athena to visualize the results of your SQL queries from the CloudFront logs.
upvoted 4 times

A1975 1 year, 4 months ago

Selected Answer: B

OptionB: Amazon Athena allows you to run standard SQL queries directly on the data stored in the S3 bucket. Amazon QuickSight is a business intelligence (BI) service that allows you to create interactive and visual dashboards to analyze data. You can connect Amazon QuickSight to Amazon Athena to visualize the results of your SQL queries from the CloudFront logs.
upvoted 2 times

ajay258 1 year, 6 months ago

Answer is B
upvoted 2 times

FFO 1 year, 8 months ago

Selected Answer: B

Athena and Quicksight. Glue is for ETL transformation
upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Answer is B
Analysis on S3 = Athena
Visualizations = Quicksight
upvoted 2 times

GalileoEC2 1 year, 8 months ago

Why the Hell A?
upvoted 1 times

GalileoEC2 1 year, 8 months ago

Why A! as far as I know Glue is not used for visualization
upvoted 2 times

Bhrino 1 year, 9 months ago

Selected Answer: B

Selected Answer: B

B because athena can be used to analyse data in s3 buckets and AWS quicksight is literally used to create visual representation of data

upvoted 2 times

LuckyAro 1 year, 9 months ago

Selected Answer: B

Using Athena to query the CloudFront logs in the S3 bucket and QuickSight to visualize the results is the best solution because it is cost-effective, scalable, and requires no infrastructure setup. It also provides a robust solution that enables the company to perform advanced analysis and build interactive visualizations without the need for a dedicated team of developers.

upvoted 2 times

skiwili 1 year, 9 months ago

Selected Answer: B

Yes B is the answer

upvoted 2 times

obatunde 1 year, 9 months ago

Selected Answer: B

Correct answer should be B.

upvoted 2 times

Namrash 1 year, 9 months ago

B is correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #281

Topic 1


A company runs a fleet of web servers using an Amazon RDS for PostgreSQL DB instance. After a routine compliance check, the company sets a standard that requires a recovery point objective (RPO) of less than 1 second for all its production databases.

Which solution meets these requirements?

- A. Enable a Multi-AZ deployment for the DB instance. **Most Voted**
- B. Enable auto scaling for the DB instance in one Availability Zone.
- C. Configure the DB instance in one Availability Zone, and create multiple read replicas in a separate Availability Zone.
- D. Configure the DB instance in one Availability Zone, and configure AWS Database Migration Service (AWS DMS) change data capture (CDC) tasks.

Correct Answer: A

Community vote distribution



Comments

KZM **Highly Voted** 1 year, 9 months ago

A:
By using Multi-AZ deployment, the company can achieve an RPO of less than 1 second because the standby instance is always in sync with the primary instance, ensuring that data changes are continuously replicated.
upvoted 24 times

A1975 **Highly Voted** 1 year, 4 months ago

Selected Answer: A

Read Replicas:
Read Replicas are asynchronous and support read scalability.
It is useful to improve performance.
Read Replicas can be in the same region or in a different region for disaster recovery purposes, but this involves manual intervention,
which means Read Replicas do not provide automatic failover and requires DNS updates and application changes

Multi-AZ:
Multi-AZ maintains a synchronous standby replica of the primary instance in a different Availability Zone within the same region.
Multi-AZ deployments provide high availability and automatic failover

multi-AZ deployments provide high availability and automatic failover.

Option A is better choice with respect to below statement,

"the company sets a standard that requires a recovery point objective (RPO) of less than 1 second for all its production databases."

upvoted 15 times

Rcosmos Most Recent 4 months, 2 weeks ago

Selected Answer: B

Solução Detalhada:

Configure o grupo de segurança do EC2 com uma regra de entrada para permitir conexões apenas do grupo de segurança do ALB.

A AWS permite essa associação dinâmica, o que garante que apenas o tráfego legítimo roteado pelo ALB possa chegar às instâncias.

Essa é uma prática recomendada e atende ao objetivo de restringir o acesso às instâncias do EC2.

upvoted 1 times

manoj6992 8 months, 1 week ago

Correct answer is A

upvoted 2 times

Uzbekistan 8 months, 1 week ago

Same question old from SAA-C02
08.2021

<https://www.examttopics.com/discussions/amazon/view/61072-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Uzbekistan 8 months, 1 week ago

Selected Answer: C

C. Configure the DB instance in one Availability Zone, and create multiple read replicas in a separate Availability Zone.

Here's why:

Read Replicas: By configuring read replicas in a separate Availability Zone, you can leverage asynchronous replication to replicate data from the primary DB instance to the read replicas with minimal latency. This setup allows for near real-time data replication and can help achieve a low RPO.

In contrast, Multi-AZ deployments (Option A) provide high availability by maintaining a standby replica but may not guarantee an RPO of less than 1 second due to synchronous replication and failover considerations. Options B and D are not directly related to achieving a low RPO for the databases.

upvoted 2 times

pentium75 11 months, 2 weeks ago

I'm unsure with A, because the term RPO is not only applied to datacenter outages. Say, an application error corrupts the database, or an administrator accidentally overwrites all records. With answer A, Multi-AZ, these changes would be instantly copied to the replica.

Only reason why A might still be correct is that the other answers don't make much more sense.

B has nothing to do with RPO at all

C could lose more than 1 second since read replicas are asynchronous

D could be part of a solution but the CDC ask alone won't help

upvoted 7 times

narddrer 1 year, 5 months ago

Selected Answer: D

option A doesn't provide Data integrity only achieved in Option D using CDC.

upvoted 1 times

FFO 1 year, 8 months ago

Selected Answer: A

Used for DR. Every single change is replicated in a standby AZ. If we lose the main AZ, (uses the same DNS name) standby becomes automatic failover and the new main DB.

upvoted 4 times

TariqKipkemei 1 year, 8 months ago

Answer is A
High availability = Multi AZ
upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: A

My vote is A
upvoted 2 times

ManOnTheMoon 1 year, 9 months ago

Agree with A
upvoted 2 times

LuckyAro 1 year, 9 months ago

Selected Answer: A

Multi-AZ is a synchronous communication with the Master in "real time" and fail over will be almost instant.
upvoted 4 times

GwonLEE 1 year, 9 months ago

Selected Answer: A

correct is A
upvoted 2 times

Namrash 1 year, 9 months ago

A should be correct
upvoted 3 times

Joan111edu 1 year, 9 months ago

Selected Answer: A

should be A
upvoted 3 times

rrharris 1 year, 9 months ago

Correct Answer is A
upvoted 8 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #282

Topic 1

A company runs a web application that is deployed on Amazon EC2 instances in the private subnet of a VPC. An Application Load Balancer (ALB) that extends across the public subnets directs web traffic to the EC2 instances. The company wants to implement new security measures to restrict inbound traffic from the ALB to the EC2 instances while preventing access from any other source inside or outside the private subnet of the EC2 instances.

Which solution will meet these requirements?

- A. Configure a route in a route table to direct traffic from the internet to the private IP addresses of the EC2 instances.
- B. Configure the security group for the EC2 instances to only allow traffic that comes from the security group for the ALB.
Most Voted
- C. Move the EC2 instances into the public subnet. Give the EC2 instances a set of Elastic IP addresses.
- D. Configure the security group for the ALB to allow any TCP traffic on any port.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Abrar2022 **Highly Voted** 1 year, 6 months ago

Read the discussion, that's the whole point why examtopics picks the wrong answer. Follow most voted answer not examtopics answer

upvoted 13 times

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: B

Configure the security group for the EC2 instances to only allow traffic that comes from the security group for the ALB

upvoted 8 times

Rcosmos **Most Recent** 4 months, 2 weeks ago

Selected Answer: B

Solução Detalhada:

Configure o grupo de segurança do EC2 com uma regra de entrada para permitir conexões apenas do grupo de segurança do

ALB.

A AWS permite essa associação dinâmica, o que garante que apenas o tráfego legítimo roteado pelo ALB possa chegar às instâncias.

Essa é uma prática recomendada e atende ao objetivo de restringir o acesso às instâncias do EC2.

upvoted 1 times

Rcosmos 4 months, 2 weeks ago

Selected Answer: B

Explicação

Requisito: Garantir que o tráfego para as instâncias do EC2 só venha do Application Load Balancer (ALB) e não de outras fontes dentro ou fora da sub-rede privada.

Como funciona:

O grupo de segurança associado às instâncias do EC2 pode ser configurado para aceitar tráfego apenas de outro grupo de segurança, neste caso, o do ALB.

Isso cria uma relação explícita, onde as instâncias do EC2 só aceitam conexões originadas pelo ALB, independentemente de sua localização (sub-rede pública ou privada).

upvoted 1 times

awslearner7 1 year, 4 months ago

can anybody explains the question?

upvoted 2 times

David_Ang 1 year, 1 month ago

is just admins fault dont worry, he just made a mistake, because "C" doesnt make any sence

upvoted 1 times

theochan 10 months, 3 weeks ago

i don't even understand what the question is trying to ask

upvoted 1 times

antropaws 1 year, 6 months ago

Selected Answer: B

It's very confusing that the system marks C as correct.

upvoted 3 times

FFO 1 year, 8 months ago

Selected Answer: B

This is B. Question already tells us they only want ONLY traffic from the ALB.

upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Answer is B

upvoted 2 times

TariqKipkemei 1 year, 2 months ago

A security group acts as a firewall that controls the traffic allowed to and from the resources in your virtual private cloud (VPC).

upvoted 2 times

GalileoEC2 1 year, 8 months ago

Why C! another cazy answer , If i am concern about security why I would want to expose my EC2 to the public internet,not make sense at all, am I

correct with this? I also go with B

upvoted 3 times

LuckyAro 1 year, 9 months ago

Selected Answer: B

B is the correct answer.

upvoted 3 times

upvoted 5 times

kpato87 1 year, 9 months ago

Selected Answer: B

configure the security group for the EC2 instances to only allow traffic that comes from the security group for the ALB. This ensures that only the traffic originating from the ALB is allowed access to the EC2 instances in the private subnet, while denying any other traffic from other sources. The other options do not provide a suitable solution to meet the stated requirements.

upvoted 5 times

Bhawesh 1 year, 9 months ago

Selected Answer: B

B. Configure the security group for the EC2 instances to only allow traffic that comes from the security group for the ALB.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #283

Topic 1

A research company runs experiments that are powered by a simulation application and a visualization application. The simulation application runs on Linux and outputs intermediate data to an NFS share every 5 minutes. The visualization application is a Windows desktop application that displays the simulation output and requires an SMB file system.


The company maintains two synchronized file systems. This strategy is causing data duplication and inefficient resource usage. The company needs to migrate the applications to AWS without making code changes to either application.

Which solution will meet these requirements?

- A. Migrate both applications to AWS Lambda. Create an Amazon S3 bucket to exchange data between the applications.
- B. Migrate both applications to Amazon Elastic Container Service (Amazon ECS). Configure Amazon FSx File Gateway for storage.
- C. Migrate the simulation application to Linux Amazon EC2 instances. Migrate the visualization application to Windows EC2 instances. Configure Amazon Simple Queue Service (Amazon SQS) to exchange data between the applications.
- D. Migrate the simulation application to Linux Amazon EC2 instances. Migrate the visualization application to Windows EC2 instances. Configure Amazon FSx for NetApp ONTAP for storage. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

LuckyAro **Highly Voted** 2 years, 3 months ago

Selected Answer: D

Amazon FSx for NetApp ONTAP provides shared storage between Linux and Windows file systems.
upvoted 28 times

JA2018 6 months, 3 weeks ago

<https://aws.amazon.com/fsx/netapp-ontap/features/#:~:text=Amazon%20FSx%20for%20NetApp%20ONTAP%20offers%20high%2Dperformance%20file%20storage,NVMe%2Dover%2DTCP%20protocols>.

upvoted 1 times

kruasan Highly Voted 2 years, 1 month ago

Selected Answer: D

This solution satisfies the needs in the following ways:

- Amazon EC2 provides a seamless migration path for the existing server-based applications without code changes. The simulation app can run on Linux EC2 instances and the visualization app on Windows EC2 instances.
- Amazon FSx for NetApp ONTAP provides highly performant file storage that is accessible via both NFS and SMB. This allows the simulation app to write to NFS shares as currently designed, and the visualization app to access the same data via SMB.
- FSx for NetApp ONTAP ensures the data is synchronized and up to date across the file systems. This addresses the data duplication issues of the current setup.
- Resources can be scaled efficiently since EC2 and FSx provide scalable compute and storage on demand.

upvoted 9 times

JA2018 6 months, 3 weeks ago

lift-and-shift aka rehost => minimal/ no code change required (caveat: it depends on the original application design).

upvoted 1 times

kruasan 2 years, 1 month ago

The other options would require more significant changes:

A. Migrating to Lambda would require re-architecting both applications and not meet the requirement to avoid code changes. S3 does not provide file system access.

B. While ECS could run the apps without code changes, FSx File Gateway only provides S3 or EFS storage, neither of which offer both NFS and SMB access. Data exchange would still be an issue.

C. Using SQS for data exchange between EC2 instances would require code changes to implement a messaging system rather than a shared file system.

upvoted 1 times

mr_kanchan 1 year, 10 months ago

How does the data duplication issue get addressed on selecting D ?

upvoted 2 times

Reckless_Jas 1 year, 9 months ago

Maybe I'm wrong, but I feel like the data is duplicated between the two types of EC2 instances. By using the FSX ONTAP will address this issue.

upvoted 2 times

pentium75 1 year, 5 months ago

Because FSx for ONTAP provides THE SAME STORAGE via NFS and SMB. Duplication issue occurred because Linux saved the files via NFS to one storage, then the data was copied to ANOTHER storage that shared it via SMB.

upvoted 3 times

JA2018 Most Recent 6 months, 3 weeks ago

Selected Answer: D

Amazon FSx for NetApp ONTAP offers high-performance file storage that's broadly accessible from Linux, Windows, and macOS compute instances via the industry-standard NFS, SMB, iSCSI, and NVMe-over-TCP protocols. It enables you to use ONTAP's widely adopted data management capabilities, like snapshots, clones, and replication, with the click of a button. In addition, it provides low-cost storage capacity that's fully elastic and virtually unlimited in size, and supports compression and deduplication to help you further reduce storage costs.

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: D

It is D but I think NetApp ONTAP is an oversell in this context. They just needed a FSx solution not a whole expensive managed service...

ABC are a lot of change to the code so D is the only choice here anyway

upvoted 3 times

djgodzilla 1 year, 5 months ago

Selected Answer: D

Amazon FSx for NetApp ONTAP

Fully managed service offering shared storage between Linux and Windows file systems (Multi & Single-AZ) up to petabytes datasets and 10+Gbps.

- Allows Multi-protocol access to data using the (NFS), Server Message Block (SMB), and Internet Small Computer Systems Interface (iSCSI) protocols
- Integrated with other AWS services (KMS, IAM, CloudTrail, amazon workspace)
- ideal solution to migrate, back up, or burst your file-based applications from on-prem to AWS without application code change.

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: D

A would require code changes

B might also require code changes, and "FSx File Gateway" provides SMB access (not NFS)

C uses SQS which has no place here

D can provide same storage via SMB and NFS

upvoted 2 times

osmk 1 year, 5 months ago

<https://docs.aws.amazon.com/fsx/latest/ONTAPGuide/what-is-fsx-ontap.html>

upvoted 2 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: D

One of the use cases for Amazon FSx for NetApp ONTAP is when you need to move workloads running on NetApp or other NFS/SMB/iSCSI servers to AWS without modifying application code or how you manage data.

upvoted 5 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

The key requirements are:

Simulation app runs on Linux, outputs data to NFS

Visualization app runs on Windows, requires SMB file system

Migrate apps to AWS without code changes

Eliminate data duplication and inefficient resource usage

upvoted 3 times

Abrar2022 2 years ago

For shared storage between Linux and windows you need to implement Amazon FSx for NetApp ONTAP

upvoted 3 times

Wael216 2 years, 3 months ago

Selected Answer: D

windows => FSX

we didn't mention containers => can't be ECS

upvoted 3 times

everfly 2 years, 3 months ago

Selected Answer: D

Amazon FSx for NetApp ONTAP is a fully managed service that provides shared file storage built on NetApp's popular ONTAP file system. It supports NFS, SMB, and iSCSI protocols² and also allows multi-protocol access to the same data

upvoted 2 times

Yechi 2 years, 3 months ago

Selected Answer: D

Amazon FSx for NetApp ONTAP is a fully-managed shared storage service built on NetApp's popular ONTAP file system. Amazon FSx for NetApp ONTAP provides the popular features, performance, and APIs of ONTAP file systems with the agility, scalability, and simplicity of a fully managed AWS service, making it easier for customers to migrate on-premises applications that rely on NAS appliances to AWS. FSx for ONTAP file systems are similar to on-premises NetApp clusters. Within each file

system that you create, you also create one or more storage virtual machines (SVMs). These are isolated file servers each with their own endpoints for NFS, SMB, and management access, as well as authentication (for both administration and end-user data access). In turn, each SVM has one or more volumes which store your data.

<https://aws.amazon.com/de/blogs/storage/getting-started-cloud-file-storage-with-amazon-fsx-for-netapp-ontap-using-netapp-management-tools/>

upvoted 5 times

zTopic 2 years, 3 months ago

Selected Answer: B

B is correct I believe

upvoted 1 times

pentium75 1 year, 5 months ago

"FSx File Gateway" is a gateway, and ECS would require code changes.

upvoted 2 times

rrharris 2 years, 3 months ago

Answer is D

upvoted 8 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #284

Topic 1

As part of budget planning, management wants a report of AWS billed items listed by user. The data will be used to create department budgets. A solutions architect needs to determine the most efficient way to obtain this report information.

Which solution meets these requirements?

- A. Run a query with Amazon Athena to generate the report.
- B. Create a report in Cost Explorer and download the report. **Most Voted**
- C. Access the bill details from the billing dashboard and download the bill.
- D. Modify a cost budget in AWS Budgets to alert with Amazon Simple Email Service (Amazon SES).

Correct Answer: B

Community vote distribution

B (100%)

Comments

DagsH **Highly Voted** 2 years, 2 months ago

Selected Answer: B

Cost Explorer looks at the usage pattern or history
upvoted 9 times

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: B

° Cost Explorer is a AWS service that allows you to view, analyze, and manage your AWS costs and usage. It provides a variety of reports that you can use to track your costs, including a report of AWS billed items listed by user.

° Creating a report in Cost Explorer is a quick and easy way to get the information you need. You can customize the report to include the specific data you want, and you can download the report in a variety of formats, including CSV, Excel, and PDF.

upvoted 7 times

Guru4Cloud 1 year, 9 months ago

This is trick question -
You need to know the differences between the billing services.

upvoted 2 times

JA2018 Most Recent 6 months, 3 weeks ago

Selected Answer: B

To save a Cost Explorer report
Open the AWS Cost Management console at <https://console.aws.amazon.com/cost-management/home>.

In the navigation pane, choose Cost Explorer Saved Reports.

Choose Create new report. This resets all of your Cost Explorer settings to your default settings.

Select a report type.

Choose Create report.

Customize your Cost Explorer settings.

Choose Save to report library.

In the Save to report library dialog box, enter a name for your report, and then choose Save report.

upvoted 1 times

JA2018 6 months, 3 weeks ago

from <https://docs.aws.amazon.com/cost-management/latest/userguide/ce-custom-reports.html>

upvoted 1 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: B

Create a report in Cost Explorer and download the report

upvoted 3 times

WheracanIstart 2 years, 2 months ago

Selected Answer: B

Cost Explorer

upvoted 2 times

pcops 2 years, 3 months ago

Answer is B

upvoted 3 times

fulingyu288 2 years, 3 months ago

Selected Answer: B

Answer is B

upvoted 4 times

rrharris 2 years, 3 months ago

Answer is B

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #285

Topic 1

A company hosts its static website by using Amazon S3. The company wants to add a contact form to its webpage. The contact form will have dynamic server-side components for users to input their name, email address, phone number, and user message. The company anticipates that there will be fewer than 100 site visits each month.

Which solution will meet these requirements MOST cost-effectively?

A. Host a dynamic contact form page in Amazon Elastic Container Service (Amazon ECS). Set up Amazon Simple Email Service (Amazon SES) to connect to any third-party email provider.


B. Create an Amazon API Gateway endpoint with an AWS Lambda backend that makes a call to Amazon Simple Email Service (Amazon SES). **Most Voted**

C. Convert the static webpage to dynamic by deploying Amazon Lightsail. Use client-side scripting to build the contact form. Integrate the form with Amazon WorkMail.

D. Create a t2.micro Amazon EC2 instance. Deploy a LAMP (Linux, Apache, MySQL, PHP/Perl/Python) stack to host the webpage. Use client-side scripting to build the contact form. Integrate the form with Amazon WorkMail.

Correct Answer: B

Community vote distribution

**Comments**

obatunde **Highly Voted** 1 year, 9 months ago

Selected Answer: B

Correct answer is B. <https://aws.amazon.com/blogs/architecture/create-dynamic-contact-forms-for-s3-static-websites-using-aws-lambda-amazon-api-gateway-and-amazon-ses/>

upvoted 14 times

kruasan **Highly Voted** 1 year, 7 months ago

Selected Answer: B

This solution is the most cost-efficient for the anticipated 100 monthly visits because:

- API Gateway charges are based on API calls. With only 100 visits, charges would be minimal.
- AWS Lambda provides compute time for the backend code in increments of 100ms, so charges would also be negligible for this workload.

- Amazon SES is used only for sending emails from the submitted contact forms. SES has a generous free tier of 62,000 emails per month, so there would be no charges for sending the contact emails.
- No EC2 instances or other infrastructure needs to be run and paid for.

upvoted 10 times

awsgeek75 Most Recent 10 months, 3 weeks ago

Option D just made me laugh!

upvoted 5 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

B is the most cost-effective solution for this use case.

The key requirements are:

Static website hosted on S3

Add a contact form with server-side processing

Low traffic website (<100 visits per month.

upvoted 4 times

rogerHS 1 year, 5 months ago

why not C

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Option C uses Lightsail which incurs charges even at low usage. Not cost effective for low traffic sites.

upvoted 3 times

pentium75 11 months, 2 weeks ago

Also it has been decided that "the contact form will have dynamic server-side components ", thus "use client-side scripting to build the contact form" is not what has been asked for.

upvoted 3 times

datz 1 year, 8 months ago

Selected Answer: B

B would be cheaper than option D,

Member only 100 site visits per month, so you are comparing API GW used 100 times a month with constantly running EC2...

upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: B

Both api gateway and lambda are serverless so charges apply only on the 100 form submissions per month

upvoted 3 times

bdp123 1 year, 9 months ago

Selected Answer: B

After looking at cost of Workmail compared to SES - probably 'B' is better

upvoted 3 times

bdp123 1 year, 9 months ago

Selected Answer: D

Create a t2 micro Amazon EC2 instance. Deploy a LAMP (Linux Apache MySQL, PHP/Perl/Python) stack to host the webpage (free open-source). Use client-side scripting to build the contact form. Integrate the form with Amazon WorkMail. This solution will provide the company with the necessary components to host the contact form page and integrate it with Amazon WorkMail at the lowest cost. Option A requires the use of Amazon ECS, which is more expensive than EC2, and Option B requires the use of Amazon API Gateway, which is also more expensive than EC2. Option C requires the use of Amazon Lightsail, which is more expensive than EC2.

<https://aws.amazon.com/what-is/lamp-stack/>

upvoted 2 times

upvoted 2 times

SkyZeroZx 1 year, 7 months ago

3 million API Gateway == 3,50 USD (EE.UU. Este (Ohio))

Is more cheaper letter B

<https://aws.amazon.com/es/api-gateway/pricing/>

<https://aws.amazon.com/es/lambda/pricing/>

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Option D uses EC2 which has a higher monthly cost than serverless options. LAMP stack adds complexity for a simple contact form.

upvoted 2 times

Palanda 1 year, 9 months ago

Selected Answer: B

It's B

upvoted 2 times

LuckyAro 1 year, 9 months ago

Selected Answer: B

B allows the company to create an API endpoint using AWS Lambda, which is a cost-effective and scalable solution for a contact form with low traffic. The backend can make a call to Amazon SES to send email notifications, which simplifies the process and reduces complexity.

upvoted 2 times

cloudbusting 1 year, 9 months ago

it is B : <https://aws.amazon.com/blogs/architecture/create-dynamic-contact-forms-for-s3-static-websites-using-aws-lambda-amazon-api-gateway-and-amazon-ses/>

upvoted 4 times

bdp123 1 year, 9 months ago

Selected Answer: D

<https://docs.aws.amazon.com/lambda/latest/dg/services-apigateway.html>

Using AWS Lambda with Amazon API Gateway - AWS Lambda

<https://docs.aws.amazon.com/lambda/latest/dg/services-apigateway.html>

<https://aws.amazon.com/lambda/faqs/>

AWS Lambda FAQs

<https://aws.amazon.com/lambda/faqs/>

upvoted 1 times

Guru4Cloud 1 year, 3 months ago

Option D uses EC2 which has a higher monthly cost than serverless options. LAMP stack adds complexity for a simple contact form.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #286

Topic 1

A company has a static website that is hosted on Amazon CloudFront in front of Amazon S3. The static website uses a database backend. The company notices that the website does not reflect updates that have been made in the website's Git repository. The company checks the continuous integration and continuous delivery (CI/CD) pipeline between the Git repository and Amazon S3. The company verifies that the webhooks are configured properly and that the CI/CD pipeline is sending messages that indicate successful deployments.


A solutions architect needs to implement a solution that displays the updates on the website.

Which solution will meet these requirements?

- A. Add an Application Load Balancer.
- B. Add Amazon ElastiCache for Redis or Memcached to the database layer of the web application.
- C. Invalidate the CloudFront cache. **Most Voted**
- D. Use AWS Certificate Manager (ACM) to validate the website's SSL certificate.

Correct Answer: C

Community vote distribution



Comments

fulingyu288 **Highly Voted** 2 years, 3 months ago

Selected Answer: C

Invalidate the CloudFront cache: The solutions architect should invalidate the CloudFront cache to ensure that the latest version of the website is being served to users.

upvoted 15 times

Damdom **Highly Voted** 1 year, 9 months ago

C. Invalidate the CloudFront cache.

Explanation:

Invalidate the CloudFront cache to ensure that the latest updates from the Git repository are reflected on the static website. When updates are made to the website's Git repository and deployed to Amazon S3, the CloudFront cache may still be serving

When updates are made to the website's content repository and deployed to Amazon S3, the CloudFront cache may still be serving the old cached content to users. By invalidating the CloudFront cache, you're instructing CloudFront to fetch fresh content from the origin (Amazon S3) and serve it to users.

upvoted 8 times

ChymKuBoy Most Recent 11 months ago

Selected Answer: C

C for sure

upvoted 2 times

jayantp04 1 year, 5 months ago

Correct C, because Invalidating the CloudFront cache will force CloudFront to fetch the latest content from Amazon S3.

Not B because not related to clear cache

upvoted 5 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: C

Invalidate the CloudFront cache so that it can read the updated static page from S3.

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

C. Invalidate the CloudFront cache

upvoted 2 times

riccardoto 1 year, 10 months ago

Selected Answer: C

C is the most reasonable cause, though the question is not well-written - "The static website uses a database backend." does not make a lot of sense to me.

upvoted 3 times

kruasan 2 years, 1 month ago

Selected Answer: B

Since the static website is hosted behind CloudFront, updates made to the S3 bucket will not be visible on the site until the CloudFront cache expires or is invalidated. By invalidating the CloudFront cache after deploying updates, the latest version in S3 will be pulled and the updates will then appear on the live site.

upvoted 1 times

RoroJ 2 years ago

Isn't that C?

upvoted 5 times

Namrash 2 years, 3 months ago

B should be the right one

upvoted 1 times

Neorem 2 years, 3 months ago

Selected Answer: C

We need to create an Cloudfront invalidation

upvoted 3 times

Bhawesh 2 years, 3 months ago

Selected Answer: C

C. Invalidate the CloudFront cache.

Problem is the CF cache. After invalidating the CloudFront cache, CF will be forced to read the updated static page from the S3 and the S3 changes will start being visible.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #287

Topic 1

A company wants to migrate a Windows-based application from on premises to the AWS Cloud. The application has three tiers: an application tier, a business tier, and a database tier with Microsoft SQL Server. The company wants to use specific features of SQL Server such as native backups and Data Quality Services. The company also needs to share files for processing between the tiers.

How should a solutions architect design the architecture to meet these requirements?

A. Host all three tiers on Amazon EC2 instances. Use Amazon FSx File Gateway for file sharing between the tiers.


B. Host all three tiers on Amazon EC2 instances. Use Amazon FSx for Windows File Server for file sharing between the tiers. **Most Voted**

C. Host the application tier and the business tier on Amazon EC2 instances. Host the database tier on Amazon RDS. Use Amazon Elastic File System (Amazon EFS) for file sharing between the tiers.

D. Host the application tier and the business tier on Amazon EC2 instances. Host the database tier on Amazon RDS. Use a Provisioned IOPS SSD (io2) Amazon Elastic Block Store (Amazon EBS) volume for file sharing between the tiers.

Correct Answer: B

Community vote distribution



Comments

KZM **Highly Voted** 2 years, 3 months ago

It is B:

A: Incorrect> FSx file Gateway designed for low latency and efficient access to in-cloud FSx for Windows File Server file shares from your on-premises facility.

B: Correct> This solution will allow the company to host all three tiers on Amazon EC2 instances while using Amazon FSx for Windows File Server to provide Windows-based file sharing between the tiers. This will allow the company to use specific features of SQL Server, such as native backups and Data Quality Services, while sharing files for processing between the tiers.

C: Incorrect> Currently, Amazon EFS supports the NFSv4.1 protocol and does not natively support the SMB protocol, and can't be used in Windows instances yet.

D: Incorrect> Amazon EBS is a block-level storage solution that is typically used to store data at the operating system level, rather than for file sharing between servers.

upvoted 19 times

djgodzilla Most Recent 1 year, 5 months ago

Selected Answer: B

RDS for SQL Backups: aren't Native MSSQL backups Instead RDS creates a storage volume snapshot of the instance, backing up the entire instance not just individual databases.

upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: B

Not A - File Gateway is just a gateway, needs S3 too

B - Yes

Not C - RDS does not support the required features, and EFS does not provide SMB

Not D - RDS does not support the required features, and EBS volume shared between tiers doesn't make sense

upvoted 4 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

B. Host all three tiers on Amazon EC2 instances. Use Amazon FSx for Windows File Server for file sharing between the tiers.

upvoted 2 times

Abrar2022 2 years ago

The question mentions Microsoft = windows

EFS is Linux

upvoted 2 times

kruasan 2 years, 1 month ago

Selected Answer: B

This design satisfies the needs in the following ways:

- Running all tiers on EC2 allows using SQL Server on EC2 with its native features like backups and Data Quality Services. SQL Server cannot be run directly on RDS.
- Amazon FSx for Windows File Server provides fully managed Windows file storage with SMB access. This allows sharing files between the Windows EC2 instances for all three tiers.
- FSx for Windows File Server has high performance, so it can handle file sharing needs between the tiers.

upvoted 2 times

kruasan 2 years, 1 month ago

The other options would not meet requirements:

A. FSx File Gateway only provides access to S3 or EFS storage. It cannot be used directly for Windows file sharing.

C. RDS cannot run SQL Server or its native tools. The database tier needs to run on EC2.

D. EBS volumes can only be attached to a single EC2 instance. They cannot be shared between tiers for file exchanges.

upvoted 2 times

Netgear 1 year, 8 months ago

No, there is RDS for SQL Server.

<https://aws.amazon.com/rds/sqlserver/>

upvoted 2 times

JA2018 6 months, 3 weeks ago

Installing SQL Server directly on an EC2 instance allows access to its native features like backups and Data Quality Services. For this use case, SQL Server cannot be run directly on RDS.

upvoted 1 times

fageroff 1 year, 7 months ago

IO2 support multi-attach

upvoted 1 times

ManOnTheMoon 2 years, 3 months ago

Whv not C?

...ing not C.

upvoted 1 times

KZM 2 years, 3 months ago

Currently, Amazon EFS supports the NFSv4.1 protocol and does not natively support the SMB protocol, and can't be used in Windows instances yet.

upvoted 4 times

AlmeroSenior 2 years, 3 months ago

Selected Answer: B

Yup B . RDS will not work , Native Backup only to S3 , and Data Quality is not supported , so all EC2 .
<https://aws.amazon.com/premiumsupport/knowledge-center/native-backup-rds-sql-server/> and <https://www.sqlserver-dba.com/2021/07/aws-rds-sql-server-limitations.html>

upvoted 3 times

LuckyAro 2 years, 3 months ago

After further research, I concur that the correct answer is B. Native Back up and Data Quality not supported on RDS for Ms SQL

upvoted 3 times

LuckyAro 2 years, 3 months ago

Selected Answer: C

C.
Host the application tier and the business tier on Amazon EC2 instances.
Host the database tier on Amazon RDS.
Use Amazon Elastic File System (Amazon EFS) for file sharing between the tiers.

This solution allows the company to use specific features of SQL Server such as native backups and Data Quality Services, by hosting the database tier on Amazon RDS. It also enables file sharing between the tiers using Amazon EFS, which is a fully managed, highly available, and scalable file system. Amazon EFS provides shared access to files across multiple instances, which is important for processing files between the tiers. Additionally, hosting the application and business tiers on Amazon EC2 instances provides the company with the flexibility to configure and manage the environment according to their requirements.

upvoted 2 times

rushi0611 2 years, 1 month ago

How are you gonna connect the EFS to windows based ??

upvoted 3 times

Yechi 2 years, 3 months ago

Selected Answer: B

Data Quality Services: If this feature is critical to your workload, consider choosing Amazon RDS Custom or Amazon EC2.
<https://docs.aws.amazon.com/prescriptive-guidance/latest/migration-sql-server/comparison.html>

upvoted 4 times

Bhawesh 2 years, 3 months ago

Selected Answer: B

Correct Answer: B

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #288

Topic 1

A company is migrating a Linux-based web server group to AWS. The web servers must access files in a shared file store for some content. The company must not make any changes to the application.

What should a solutions architect do to meet these requirements?

- A. Create an Amazon S3 Standard bucket with access to the web servers.
- B. Configure an Amazon CloudFront distribution with an Amazon S3 bucket as the origin.
- C. Create an Amazon Elastic File System (Amazon EFS) file system. Mount the EFS file system on all web servers.
- D. Configure a General Purpose SSD (gp3) Amazon Elastic Block Store (Amazon EBS) volume. Mount the EBS volume to all web servers.

Most Voted

Correct Answer: C

Community vote distribution

C (100%)

Comments

Bhawesh Highly Voted 2 years, 3 months ago

Selected Answer: C

Since no code change is permitted, below choice makes sense for the unix server's file sharing:
C. Create an Amazon Elastic File System (Amazon EFS) file system. Mount the EFS file system on all web servers.

upvoted 20 times

TariqKipkemei Most Recent 1 year, 8 months ago

Selected Answer: C

Rehost the application web servers on EC2 and Create an Amazon Elastic File System (Amazon EFS) file system. Mount the EFS file system on all web servers.

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

Since no code change is permitted, below choice makes sense for the unix server's file sharing:

upvoted 3 times

callmejaja 1 year, 11 months ago

Selected Answer: C

Since no code change is permitted, below choice makes sense for the unix server's file sharing:

upvoted 3 times

antropaws 2 years ago

Selected Answer: C

C is correct.

upvoted 2 times

kruasan 2 years, 1 month ago

Selected Answer: C

This solution satisfies the needs in the following ways:

- EFS provides a fully managed elastic network file system that can be mounted on multiple EC2 instances concurrently.
- The EFS file system appears as a standard file system mount on the Linux web servers, requiring no application changes. The servers can access shared files as if they were on local storage.
- EFS is highly available, durable, and scalable, providing a robust shared storage solution.

upvoted 3 times

kruasan 2 years, 1 month ago

The other options would require modifying the application or do not provide a standard file system:

A. S3 does not provide a standard file system mount or share. The application would need to be changed to access S3 storage.

B. CloudFront is a content delivery network and caching service. It does not provide a file system mount or share and would require application changes.

D. EBS volumes can only attach to a single EC2 instance. They cannot be mounted by multiple servers concurrently and do not provide a shared file system.

upvoted 4 times

Steve_4542636 2 years, 3 months ago

Selected Answer: C

No application changes are allowed and EFS is compatible with Linux

upvoted 2 times

LuckyAro 2 years, 3 months ago

Selected Answer: C

C is the answer:

Create an Amazon Elastic File System (Amazon EFS) file system.

Mount the EFS file system on all web servers.

To meet the requirements of providing a shared file store for Linux-based web servers without making changes to the application, using an Amazon EFS file system is the best solution.

Amazon EFS is a managed NFS file system service that provides shared access to files across multiple Linux-based instances, which makes it suitable for this use case.

Amazon S3 is not ideal for this scenario since it is an object storage service and not a file system, and it requires additional tools or libraries to mount the S3 bucket as a file system.

Amazon CloudFront can be used to improve content delivery performance but is not necessary for this requirement.

Additionally, Amazon EBS volumes can only be mounted to one instance at a time, so it is not suitable for sharing files across multiple instances.

upvoted 3 times

Karlos99 2 years, 3 months ago

But what about aws ebs multi attach?

upvoted 2 times

JA2018 6 months, 3 weeks ago

Key in STEM:

The web servers must access files in a shared file store for some content

upvoted 1 times

elearningtakai 2 years, 2 months ago

Amazon EBS Multi-Attach enables you to attach a single Provisioned IOPS SSD (io1 or io2) volume to multiple instances. EBS General Purpose SSD (gp3) doesn't support Multi-Attach

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #289

Topic 1

A company has an AWS Lambda function that needs read access to an Amazon S3 bucket that is located in the same AWS account.

Which solution will meet these requirements in the MOST secure manner?

A. Apply an S3 bucket policy that grants read access to the S3 bucket.

B. Apply an IAM role to the Lambda function. Apply an IAM policy to the role to grant read access to the S3 bucket.

Most Voted

C. Embed an access key and a secret key in the Lambda function's code to grant the required IAM permissions for read access to the S3 bucket.

D. Apply an IAM role to the Lambda function. Apply an IAM policy to the role to grant read access to all S3 buckets in the account.

Correct Answer: B

Community vote distribution

B (100%)

Comments

kruasan Highly Voted 1 year, 7 months ago

Selected Answer: B

This solution satisfies the needs in the most secure manner:

- An IAM role provides temporary credentials to the Lambda function to access AWS resources. The function does not have persistent credentials.
- The IAM policy grants least privilege access by specifying read access only to the specific S3 bucket needed. Access is not granted to all S3 buckets.
- If the Lambda function is compromised, the attacker would only gain access to the one specified S3 bucket. They would not receive broad access to resources.

upvoted 9 times

kruasan 1 year, 7 months ago

The other options are less secure:

A. A bucket policy grants open access to a resource. It is a less granular way to provide access and grants more privilege than needed.

C. Embedding access keys in code is extremely insecure and against best practices. The keys provide full access and are at major risk of compromise if the code leaks.
D. Granting access to all S3 buckets provides far too much privilege if only one bucket needs access. It greatly expands the impact if compromised.

upvoted 4 times

Rido4good Most Recent 10 months, 3 weeks ago

Has anyone passed this exam, choosing the wrong answers from ExamTopics? or what's the reason for the confusion???

upvoted 1 times

bbgun891404021 11 months, 1 week ago

Selected Answer: B

B is correct.

upvoted 3 times

TMabs 1 year, 2 months ago

Answer=B

upvoted 2 times

antropaws 1 year, 6 months ago

Selected Answer: B

B is correct.

upvoted 2 times

Dr_Chomp 1 year, 8 months ago

Selected Answer: B

you dont want to grant access to all S3 buckets (which is answer D) - only the one identified (so answer A)

upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: B

B is only for one bucket and you want to use Role based security here.

upvoted 2 times

Ja13 1 year, 9 months ago

Selected Answer: B

C, it says MOST secure manner, so only to one bucket

upvoted 1 times

Joxtat 1 year, 9 months ago

Selected Answer: B

<https://docs.aws.amazon.com/lambda/latest/dg/lambda-permissions.html>

upvoted 2 times

kpato87 1 year, 9 months ago

Selected Answer: B

This is the most secure and recommended way to provide an AWS Lambda function with access to an S3 bucket. It involves creating an IAM role that the Lambda function assumes, and attaching an IAM policy to the role that grants the necessary permissions to read from the S3 bucket.

upvoted 4 times

Joan111edu 1 year, 9 months ago

Selected Answer: B

B. Least of privilege

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #290

Topic 1


A company hosts a web application on multiple Amazon EC2 instances. The EC2 instances are in an Auto Scaling group that scales in response to user demand. The company wants to optimize cost savings without making a long-term commitment.

Which EC2 instance purchasing option should a solutions architect recommend to meet these requirements?

- A. Dedicated Instances only
- B. On-Demand Instances only
- C. A mix of On-Demand Instances and Spot Instances **Most Voted**
- D. A mix of On-Demand Instances and Reserved Instances

Correct Answer: C

Community vote distribution



Comments

Mikado211 **Highly Voted** 1 year, 6 months ago

Selected Answer: C

There is a little trap here, because in the way this question is asked, both B and C are true since we don't know if it's production or not.

In a production environment C is absolutely forbidden and B is the good solution, we even have questions in this dump about this case.

In a dev environment spot instances are good if you don't care about stability so C can be a good answer.

Since this question is all about cost, let's go for the stingy rat solution, spot instances are cheaper, so C is correct.

upvoted 14 times

Uzbekistan **Highly Voted** 1 year, 2 months ago

Selected Answer: C

By using a mix of On-Demand Instances and Spot Instances, the company can leverage the cost-effectiveness of Spot Instances for parts of their workload while ensuring the availability and reliability of On-Demand Instances for critical components. This approach allows for cost optimization without sacrificing performance or reliability.

upvoted 8 times

ExamGuru727 Most Recent 1 year, 2 months ago

Selected Answer: C

On-demand + spot for additional capacity will save costs.
upvoted 3 times

thewalker 1 year, 4 months ago

Selected Answer: B

On-Demand.
upvoted 1 times

Krishhhh 1 year, 6 months ago

Any one have dumps
upvoted 1 times

beginnercloud 1 year, 7 months ago

Selected Answer: C

It's about COST, not operational efficiency for this question :) C is correct
upvoted 5 times

TariqKipkemei 1 year, 8 months ago

Selected Answer: C

A mix of On-Demand Instances to handle baseline workload and Spot Instances to handle excess workload.
upvoted 5 times

Kt 1 year, 8 months ago

Exam topic is not free anymore. Anyone has free access ?
upvoted 1 times

Kaula 1 year, 2 months ago

Free up to question 400, from 401 you have to pay
upvoted 1 times

AWSSURI 9 months, 2 weeks ago

Now its upto question 250 and im pretty sure one day they will ask 100 dollars for all of 900 questions
upvoted 2 times

soewailin 1 year, 8 months ago

for now though, I have still access.
upvoted 3 times

Damdom 1 year, 9 months ago

Selected Answer: C

By combining On-Demand Instances for steady-state workloads or critical components and Spot Instances for less critical or burstable workloads, you can achieve a balance between cost savings and performance. This strategy allows you to optimize costs without making a long-term commitment, as Spot Instances provide cost savings without the need for upfront payments or long-term contracts.
upvoted 4 times

Abrar2022 2 years ago

Selected Answer: C

It's about COST, not operational efficiency for this question.
upvoted 4 times

kraken21 2 years, 2 months ago

Selected Answer: C

Autoscaling with ALB / scale up on demand using on demand and spot instance combination makes sense. Reserved will not fit the no-long term commitment clause.

upvoted 2 times

WheracanIstart 2 years, 2 months ago

Selected Answer: C

Without commitment....Spot instances

upvoted 2 times

cegama543 2 years, 2 months ago

Selected Answer: B

If the company wants to optimize cost savings without making a long-term commitment, then using only On-Demand Instances may not be the most cost-effective option. Spot Instances can be significantly cheaper than On-Demand Instances, but they come with the risk of being interrupted if the Spot price increases above your bid price. If the company is willing to accept this risk, a mix of On-Demand Instances and Spot Instances may be the best option to optimize cost savings while maintaining the desired level of scalability.

However, if the company wants the most predictable pricing and does not want to risk instance interruption, then using only On-Demand Instances is a good choice. It ultimately depends on the company's priorities and risk tolerance.

upvoted 3 times

pentium75 1 year, 5 months ago

First, the question is about cost, cost, cost. Second, answer C is "a mix of on-demand and spot instances"; they could still use on-demand if spot is not available.

upvoted 3 times

Steve_4542636 2 years, 3 months ago

Selected Answer: C

It's about COST, not operational efficiency for this question.

upvoted 2 times

Samuel03 2 years, 3 months ago

Selected Answer: C

Should be C

upvoted 2 times

bdp123 2 years, 3 months ago

Selected Answer: C

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/ec2-auto-scaling-mixed-instances-groups.html>

upvoted 2 times

AlmeroSenior 2 years, 3 months ago

Selected Answer: C

C - WEB apps , mostly Stateless , and ASG support OnDemand and Spot mix , in fact , you can prioritize to have Ondemand , before it uses Spot > <https://docs.aws.amazon.com/autoscaling/ec2/userguide/launch-template-spot-instances.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #291

Topic 1

A media company uses Amazon CloudFront for its publicly available streaming video content. The company wants to secure the video content that is hosted in Amazon S3 by controlling who has access. Some of the company's users are using a custom HTTP client that does not support cookies. Some of the company's users are unable to change the hardcoded URLs that they are using for access.

Which services or methods will meet these requirements with the LEAST impact to the users? (Choose two.)

A. Signed cookies **Most Voted**

B. Signed URLs **Most Voted**


C. AWS AppSync

D. JSON Web Token (JWT)

E. AWS Secrets Manager

Correct Answer: AB

Community vote distribution



Comments

leoattf **Highly Voted** 1 year, 9 months ago

Selected Answer: AB

I thought that option A was totally wrong, because the question mentions "HTTP client does not support cookies". However it is right, along with option B. Check the link bellow, first paragraph.

<https://aws.amazon.com/blogs/media/secure-content-using-cloudfront-functions/>

upvoted 29 times

Steve_4542636 1 year, 9 months ago

Thanks for this! What a tricky question. If the client doesn't support cookies, THEN they use the signed S3 Urls.

upvoted 11 times

AAAWrekng 1 year, 1 month ago

LOL, like the old question, in my hand I have 2 coins, and they equal 15 cents, one of them is not a nickel. What are the coins

upvoted 6 times

bujuman 9 months ago

Plus, Customers can choose to use either one or both, depending on the use case.
Thanks for this share !

upvoted 3 times

johnmcclane78 Highly Voted 1 year, 9 months ago

B. Signed URLs - This method allows the media company to control who can access the video content by creating a time-limited URL with a cryptographic signature. This URL can be distributed to the users who are unable to change the hardcoded URLs they are using for access, and they can access the content without needing to support cookies.

D. JSON Web Token (JWT) - This method allows the media company to control who can access the video content by creating a secure token that contains user authentication and authorization information. This token can be distributed to the users who are using a custom HTTP client that does not support cookies. The users can include this token in their requests to access the content without needing to support cookies.

Therefore, options B and D are the correct answers.

Option A (Signed cookies) would not work for users who are using a custom HTTP client that does not support cookies. Option C (AWS AppSync) is not relevant to the requirement of securing video content. Option E (AWS Secrets Manager) is a service used for storing and retrieving secrets, which is not relevant to the requirement of securing video content.

upvoted 21 times

pentium75 11 months, 2 weeks ago

So you want to 'distribute the signed URL to the users who are unable to change the hardcoded URL'?

upvoted 2 times

ONS_KH 1 year, 1 month ago

This is the response of chatgpt isnt it ? Pay attention ! it doesn't always give the right answer

upvoted 5 times

Dantecito Most Recent 4 months ago

Selected Answer: AB

Signed cookies, signed URLs, and JSON web token (JWT) are Amazon Cloudfront authentication methods. But as the question says with the least possible impact for users, I would go with AB since D requires using Lambda@Edge and I suppose it has more impact to the users than the first two.

upvoted 1 times

Uzbekistan 8 months, 1 week ago

Selected Answer: BD

B. Signed URLs: Signed URLs allow you to control access to your content in CloudFront by providing URLs that are valid only for a specified duration. This means users can access the content using the same URLs they have hardcoded, without the need for cookies or special client support.

D. JSON Web Token (JWT): JSON Web Tokens (JWTs) can be used to control access to resources by embedding authentication and authorization information in the token itself. Users can include the JWT in the request headers, allowing access to be controlled without relying on cookies. This approach doesn't require changes to hardcoded URLs and can be integrated into custom HTTP clients.

upvoted 1 times

lostmagnet001 10 months ago

Selected Answer: AB

a little tricky but you have to "control" access, ok dont support cookies, so put signed cookies.

upvoted 2 times

ray320x 10 months, 1 week ago

so how many marks do you get if you get 1 wrong

upvoted 5 times

awsgeek75 10 months, 3 weeks ago

Selected Answer: AB

If you have gotten this far and got THIS trick question right then you are going to make it! Good Luck!

upvoted 6 times

pentium75 11 months, 2 weeks ago

Selected Answer: AB

'SOME are using a client that does not support cookies' -> use signed URLs

'SOME are unable to change the hardcoded URLs' -> used signed cookies

upvoted 16 times

ale_brd_111 11 months, 2 weeks ago

Selected Answer: AB

Signed URLs and signed cookies are the most suitable options. They can effectively address the requirements of both users with custom HTTP clients and those with hardcoded URLs.

upvoted 2 times

prabhjot 1 year, 2 months ago

B & E - B. Signed URLs: This allows you to generate time-limited URLs with a signature that grants temporary access to specific resources in your S3 bucket. It doesn't rely on cookies and can be generated for users without requiring any changes to their HTTP client or hardcoded URLs. This method provides fine-grained control over access to your content.

E. AWS Secrets Manager: While AWS Secrets Manager can be useful for managing and rotating secrets, it is not directly related to securing S3 content in the context of the question. It's not one of the primary methods for securing access to S3 objects.

upvoted 1 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: AB

To secure streaming video content from Amazon CloudFront, two methods are available: signed cookies or signed URLs. Customers can choose to use either one or both, depending on the use case.

upvoted 4 times

tabbyDolly 1 year, 2 months ago

AB - <https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/PrivateContent.html>

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: BD

B and D are the correct options for meeting the requirements with the least impact to users.

Signed URLs allow access to individual objects in Amazon S3 for a specified time period without requiring cookies. This allows the custom HTTP client users to access content.

JSON Web Tokens (JWT) allow users to get temporary access tokens that can be passed in requests. This allows users with hardcoded URLs to access content without updating URLs.

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

No good

Signed cookies require client support and may impact users.

AWS AppSync and Secrets Manager do not help address the specific access requirements.

Good

So Signed URLs and JWTs allow securing access to S3 content with minimal impact to users, meeting the requirements.

upvoted 2 times

riccardoto 1 year, 4 months ago

Selected Answer: BD

I understand why many users here are voting AB, but in my opinion BD is more correct

I understand why many users here are voting AB, but in my opinion BC is more correct.

Using JWT or signed urls will work both for users that cannot use cookies or cannot change the url.

upvoted 2 times

katetel 1 year, 4 months ago

Selected Answer: AB

it's correct

upvoted 2 times

MrAWSAssociate 1 year, 5 months ago

Selected Answer: CE

These are the right answers!

upvoted 2 times

DrWatson 1 year, 6 months ago

Selected Answer: AB

"Some of the company's users" does not support cookies, then they'll use Signed URLs.

"Some of the company's users" are unable to change the hardcoded URLs, then they'll use Signed cookies.

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #292

Topic 1

A company is preparing a new data platform that will ingest real-time streaming data from multiple sources. The company needs to transform the data before writing the data to Amazon S3. The company needs the ability to use SQL to query the transformed data.

Which solutions will meet these requirements? (Choose two.)

A. Use Amazon Kinesis Data Streams to stream the data. Use Amazon Kinesis Data Analytics to transform the data. Use Amazon Kinesis Data Firehose to write the data to Amazon S3. Use Amazon Athena to query the transformed data from Amazon S3. **Most Voted**

B. Use Amazon Managed Streaming for Apache Kafka (Amazon MSK) to stream the data. Use AWS Glue to transform the data and to write the data to Amazon S3. Use Amazon Athena to query the transformed data from Amazon S3. **Most Voted**

C. Use AWS Database Migration Service (AWS DMS) to ingest the data. Use Amazon EMR to transform the data and to write the data to Amazon S3. Use Amazon Athena to query the transformed data from Amazon S3.

D. Use Amazon Managed Streaming for Apache Kafka (Amazon MSK) to stream the data. Use Amazon Kinesis Data Analytics to transform the data and to write the data to Amazon S3. Use the Amazon RDS query editor to query the transformed data from Amazon S3.

E. Use Amazon Kinesis Data Streams to stream the data. Use AWS Glue to transform the data. Use Amazon Kinesis Data Firehose to write the data to Amazon S3. Use the Amazon RDS query editor to query the transformed data from Amazon S3.

Correct Answer: AB

Community vote distribution

AB (93%)

AE (7%)

Comments

Steve_4542636 **Highly Voted** 1 year, 9 months ago

Selected Answer: AB

OK, for B I did some research, <https://docs.aws.amazon.com/glue/latest/dg/add-job-streaming.html>

"You can create streaming extract, transform, and load (ETL) jobs that run continuously, consume data from streaming sources like Amazon Kinesis Data Streams, Apache Kafka, and Amazon Managed Streaming for Apache Kafka (Amazon MSK). The jobs cleanse and transform the data, and then load the results into Amazon S3 data lakes or JDBC data stores."

upvoted 16 times

Paras043 Highly Voted 1 year, 8 months ago

But how can you transform data using kinesis data analytics ??

upvoted 8 times

luisgu 1 year, 7 months ago

See <https://aws.amazon.com/kinesis/data-analytics/faqs/?nc=sn&loc=6>

upvoted 3 times

awsgeek75 Most Recent 10 months, 3 weeks ago

Selected Answer: AB

Just because C is not going to work a DE use RDS so totally illogical

A & B seem to have redundant streaming, transformation and query steps so not sure if these are the right choices but CDE are completely wrong anyway!

upvoted 5 times

farnamjam 11 months, 2 weeks ago

For A didn't know that Kinesis Analytics can transform the data as well:

Amazon Kinesis Data Analytics provides built-in functions to filter, aggregate, and transform streaming data for advanced analytics. It processes streaming data with sub-second latencies, enabling you to analyze and respond to incoming data and streaming events in real time.

upvoted 5 times

pentium75 11 months, 2 weeks ago

Selected Answer: AB

"Use SQL to query the transformed data" [which is in S3] requires Athena, thus D and E are out. DMS is nonsense here thus C is out.

upvoted 3 times

MiniYang 1 year ago

why E is not right choice?

upvoted 2 times

pentium75 11 months, 2 weeks ago

Because you can't "use the Amazon RDS query editor to query .. data from S3"

upvoted 2 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: AB

options A and B will meet these requirements.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: AB

A and B are correct.

A uses Kinesis Data Streams for streaming, Kinesis Data Analytics for transformation, Kinesis Data Firehose for writing to S3, and Athena for SQL queries on S3 data.

B uses Amazon MSK for streaming, AWS Glue for transformation and writing to S3, and Athena for SQL queries on S3 data.

upvoted 3 times

Diqian 1 year, 3 months ago

Why E is incorrect?

upvoted 3 times

awsgeek75 10 months, 3 weeks ago

"Use the Amazon RDS query editor to query the transformed data from Amazon S3." is not possible as RDS query editor is for RDS and not for S3

upvoted 3 times

MrCloudy 1 year, 7 months ago

Selected Answer: AE

To transform real-time streaming data from multiple sources, write it to Amazon S3, and query the transformed data using SQL, the company can use the following solutions: Amazon Kinesis Data Streams, Amazon Kinesis Data Analytics, and Amazon Kinesis Data Firehose. The transformed data can be queried using Amazon Athena. Therefore, options A and E are the correct answers.

Option A is correct because it uses Amazon Kinesis Data Streams to stream data from multiple sources, Amazon Kinesis Data Analytics to transform the data, and Amazon Kinesis Data Firehose to write the data to Amazon S3. Amazon Athena can be used to query the transformed data in Amazon S3.

Option E is also correct because it uses Amazon Kinesis Data Streams to stream data from multiple sources, AWS Glue to transform the data, and Amazon Kinesis Data Firehose to write the data to Amazon S3. Amazon Athena can be used to query the transformed data in Amazon S3.

upvoted 3 times

sand444 1 year, 2 months ago

Amazon Athena is not in option E

upvoted 8 times

kraken21 1 year, 8 months ago

Selected Answer: AB

DMS can move data from DBs to streaming services and cannot natively handle streaming data. Hence A.B makes sense. Also AWS Glue/ETL can handle MSK streaming <https://docs.aws.amazon.com/glue/latest/dg/add-job-streaming.html>.

upvoted 3 times

elearningtakai 1 year, 8 months ago

Selected Answer: AB

The solutions that meet the requirements of streaming real-time data, transforming the data before writing to S3, and querying the transformed data using SQL are A and B.

Option C: This option is not ideal for streaming real-time data as AWS DMS is not optimized for real-time data ingestion.

Option D & E: These options are not recommended as the Amazon RDS query editor is not designed for querying data in S3, and it is not efficient for running complex queries.

upvoted 6 times

gold4otas 1 year, 8 months ago

Selected Answer: AB

The correct answers are options A & B

upvoted 2 times

TungPham 1 year, 9 months ago

may Amazon RDS query editor to query the transformed data from Amazon S3 ?

i don't think so, plz get link docs to that

upvoted 1 times

ManOnTheMoon 1 year, 9 months ago

Why not A & D?

upvoted 1 times

TungPham 1 year, 9 months ago

may Amazon RDS query editor to query the transformed data from Amazon S3 ?

i don't think so, plz get link docs to that

upvoted 1 times

LuckyAro 1 year, 9 months ago

Selected Answer: AB

A and B

upvoted 2 times

designmood22 1 year, 9 months ago

Answer is : A & B

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #293

Topic 1

A company has an on-premises volume backup solution that has reached its end of life. The company wants to use AWS as part of a new backup solution and wants to maintain local access to all the data while it is backed up on AWS. The company wants to ensure that the data backed up on AWS is automatically and securely transferred.

Which solution meets these requirements?

- A. Use AWS Snowball to migrate data out of the on-premises solution to Amazon S3. Configure on-premises systems to mount the Snowball S3 endpoint to provide local access to the data.
- B. Use AWS Snowball Edge to migrate data out of the on-premises solution to Amazon S3. Use the Snowball Edge file interface to provide on-premises systems with local access to the data.
- C. Use AWS Storage Gateway and configure a cached volume gateway. Run the Storage Gateway software appliance on premises and configure a percentage of data to cache locally. Mount the gateway storage volumes to provide local access to the data.
- D. Use AWS Storage Gateway and configure a stored volume gateway. Run the Storage Gateway software appliance on premises and map the gateway storage volumes to on-premises storage. Mount the gateway storage volumes to provide local access to the data. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Steve_4542636 **Highly Voted** 1 year, 3 months ago

Selected Answer: D

The question states, "wants to maintain local access to all the data" This is storage gateway. Cached gateway stores only the frequently accessed data locally which is not what the problem statement asks for.

upvoted 21 times

kruasan **Highly Voted** 1 year, 1 month ago

Selected Answer: D

1. The company wants to maintain local access to all the data. Only stored volumes keep the complete dataset on-premises,

providing low-latency access. Cached volumes only cache a subset locally.

2. The company wants the data backed up on AWS. With stored volumes, periodic backups (snapshots) of the on-premises data are sent to S3, providing durable and scalable backup storage.

3. The company wants the data transfer to AWS to be automatic and secure. Storage Gateway provides an encrypted connection between the on-premises gateway and AWS storage. Backups to S3 are sent asynchronously and automatically based on the backup schedule configured.

upvoted 14 times

AdamVigas Most Recent 3 months, 3 weeks ago

Selected Answer: D

local access to all data

upvoted 1 times

TariqKipkemei 8 months, 1 week ago

Selected Answer: D

The Volume Gateway runs in either a cached or stored mode.

In the cached mode, your primary data is written to S3, while retaining your frequently accessed data locally in a cache for low-latency access.

In the stored mode, your primary data is stored locally and your entire dataset is available for low-latency access while asynchronously backed up to AWS.

<https://aws.amazon.com/storagegateway/faqs/#:~:text=What%20is%20Volume%20Gateway%3F>

upvoted 9 times

Guru4Cloud 9 months ago

Selected Answer: D

@kruasan well explained

upvoted 3 times

ChrisG1454 1 year, 3 months ago

Ans = D

<https://docs.aws.amazon.com/storagegateway/latest/vgw/WhatIsStorageGateway.html>

upvoted 4 times

Neha999 1 year, 3 months ago

D

<https://www.examttopics.com/discussions/amazon/view/43725-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 3 times

bdp123 1 year, 3 months ago

Selected Answer: D

<https://aws.amazon.com/storagegateway/faqs/#:~:text=In%20the%20cached%20mode%2C%20your,asynchronously%20backed%20up%20to%20AWS.>

In the cached mode, your primary data is written to S3, while retaining your frequently accessed data locally in a cache for low-latency access.

In the stored mode, your primary data is stored locally and your entire dataset is available for low-latency access while asynchronously backed up to AWS.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #294

Topic 1

An application that is hosted on Amazon EC2 instances needs to access an Amazon S3 bucket. Traffic must not traverse the internet.

How should a solutions architect configure access to meet these requirements?

- A. Create a private hosted zone by using Amazon Route 53.
- B. Set up a gateway VPC endpoint for Amazon S3 in the VPC. **Most Voted**
- C. Configure the EC2 instances to use a NAT gateway to access the S3 bucket.
- D. Establish an AWS Site-to-Site VPN connection between the VPC and the S3 bucket.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Guru4Cloud **Highly Voted** 9 months ago

Selected Answer: B

The correct answer is B. Set up a gateway VPC endpoint for Amazon S3 in the VPC.

A gateway VPC endpoint is a private way for Amazon EC2 instances in a VPC to access AWS services, such as Amazon S3, without having to go through the internet. This can help to improve security and performance.

upvoted 8 times

Steve_4542636 **Highly Voted** 1 year, 3 months ago

Selected Answer: B

S3 and DynamoDB are the only services with Gateway endpoint options

upvoted 7 times

TariqKipkemei **Most Recent** 8 months, 1 week ago

Selected Answer: B

Set up a gateway VPC endpoint for Amazon S3 in the VPC.

upvoted 4 times

ManOnTheMoon 1 year, 3 months ago

Agree with B
upvoted 2 times

jennyka76 1 year, 3 months ago

ANSWER - B
<https://docs.aws.amazon.com/vpc/latest/privatelink/gateway-endpoints.html> R B
upvoted 2 times

LuckyAro 1 year, 3 months ago

Selected Answer: B

B is correct
upvoted 2 times

skiwili 1 year, 3 months ago

Selected Answer: B

Bbbbbbbb
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #295

Topic 1


An ecommerce company stores terabytes of customer data in the AWS Cloud. The data contains personally identifiable information (PII). The company wants to use the data in three applications. Only one of the applications needs to process the PII. The PII must be removed before the other two applications process the data.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Store the data in an Amazon DynamoDB table. Create a proxy application layer to intercept and process the data that each application requests.
- B. Store the data in an Amazon S3 bucket. Process and transform the data by using S3 Object Lambda before returning the data to the requesting application. **Most Voted**
- C. Process the data and store the transformed data in three separate Amazon S3 buckets so that each application has its own custom dataset. Point each application to its respective S3 bucket.
- D. Process the data and store the transformed data in three separate Amazon DynamoDB tables so that each application has its own custom dataset. Point each application to its respective DynamoDB table.

Correct Answer: B

Community vote distribution



Comments

Steve_4542636 **Highly Voted** 2 years, 3 months ago

Selected Answer: B

Actually this is what Macie is best used for.
upvoted 19 times

1e22522 10 months, 1 week ago

"Akshually" bro go to sleep
upvoted 2 times

Mikado211 1 year, 6 months ago

Yes. That's the problem here, Macie is the recommended tool in such case, but you do not have it in the answers.

upvoted 3 times

awsgeek75 1 year, 4 months ago

Macie is for identifying the PII data. Here it's much simpler because one of the apps need the PII data and other apps don't so you don't need to identify the PII data as you know it is already there. You just need to identify the app that needs the data which is not the best use case for Macie

upvoted 4 times

fruto123 **Highly Voted** 2 years, 3 months ago

Selected Answer: B

B is the right answer and the proof is in this link.

<https://aws.amazon.com/blogs/aws/introducing-amazon-s3-object-lambda-use-your-code-to-process-data-as-it-is-being-retrieved-from-s3/>

upvoted 12 times

Guru4Cloud 1 year, 9 months ago

This is so wrong

upvoted 1 times

pentium75 1 year, 5 months ago

But it matches the exact use case here.

upvoted 2 times

awsgeek75 1 year, 4 months ago

Why do you think this is wrong?

upvoted 2 times

Rcosmos **Most Recent** 4 months, 2 weeks ago

Selected Answer: B

Por que esta é a melhor solução?

Processamento sob demanda:

O S3 Object Lambda permite que você processe e transforme objetos armazenados no S3 dinamicamente, com base na solicitação do aplicativo.

Para os dois aplicativos que não precisam das PII, o Object Lambda pode remover essas informações no momento da solicitação, sem duplicar ou alterar os dados originais.

upvoted 1 times

AWSSURI 9 months, 2 weeks ago

I miss you Burugudystunstuguy....Man's been consistent with his answers

upvoted 2 times

MikeJANG 1 year, 3 months ago

Selected Answer: C

[GPT4] while S3 Object Lambda is a powerful tool for real-time data transformation, it is not the best fit for processing very large datasets due to Lambda's execution limits(15 min). Instead, preprocessing the data and storing it in separate S3 buckets for each application's needs is a more operationally efficient solution for the scenario describes.

upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: B

Because this is exactly what the AWS blog says.

"When you store data in Amazon Simple Storage Service (Amazon S3), you can easily share it for use by multiple applications. However, each application has its own requirements and may need a different view of the data. For example, a dataset created by an e-commerce application may include personally identifiable information (PII) that is not needed when the same data is processed for analytics and should be redacted."

upvoted 7 times

pentium75 1 year, 5 months ago

"Today, I'm very happy to announce the availability of S3 Object Lambda, a new capability that allows you to add your own code to process data retrieved from S3 before returning it to an application. S3 Object Lambda works with your existing applications and uses AWS Lambda functions to automatically process and transform your data as it is being retrieved from S3. The Lambda function is invoked inline with a standard S3 GET request, so you don't need to change your application code."

<https://aws.amazon.com/de/blogs/aws/introducing-amazon-s3-object-lambda-use-your-code-to-process-data-as-it-is-being-retrieved-from-s3/>

upvoted 8 times

meowruki 1 year, 6 months ago

B. Store the data in an Amazon S3 bucket. Process and transform the data by using S3 Object Lambda before returning the data to the requesting application.

This solution allows you to use S3 Object Lambda to process and transform the data on-the-fly as it is requested by each application. S3 Object Lambda enables you to apply custom code to your data retrieval requests, allowing you to remove PII before returning the data to the requesting application. This eliminates the need to create and manage separate storage locations for each application, reducing operational overhead.

upvoted 3 times

rvca231 1 year, 7 months ago

Selected Answer: C

Why would you reprocess the data every time you request it when you can just filter it once and be done?

Because of this I think A and B are highly inefficient, leaving us with C and D as options.

Since S3 is better suited for Data Lakes, I think C is the answer.

upvoted 2 times

pentium75 1 year, 5 months ago

Because this is exactly what the AWS blog says.

"When you store data in Amazon Simple Storage Service (Amazon S3), you can easily share it for use by multiple applications. However, each application has its own requirements and may need a different view of the data. For example, a dataset created by an e-commerce application may include personally identifiable information (PII) that is not needed when the same data is processed for analytics and should be redacted. ... Today, I'm very happy to announce the availability of S3 Object Lambda, a new capability that allows you to add your own code to process data retrieved from S3 before returning it to an application. S3 Object Lambda works with your existing applications and uses AWS Lambda functions to automatically process and transform your data as it is being retrieved from S3."

upvoted 2 times

awsgeek75 1 year, 4 months ago

Least operational overhead. The DevOps team is throwing this problem to the developers which is why C is not best.

upvoted 2 times

Abrar2022 2 years ago

Selected Answer: B

Store the data in an Amazon S3 bucket and using S3 Object Lambda to process and transform the data before returning it to the requesting application. This approach allows the PII to be removed in real-time and without the need to create separate datasets or tables for each application.

upvoted 2 times

antropaws 2 years ago

Selected Answer: A

@fruto123 and everyone that upvoted:

Is it plausible that S3 Object Lambda can process terabytes of data in 60 seconds? The same link you shared states that the maximum duration for a Lambda function used by S3 Object Lambda is 60 seconds.

Answer is A.

upvoted 2 times

antropaws 2 years ago

Chat GPT:

Isn't just 60 seconds the maximum duration for a Lambda function used by S3 Object Lambda? How can it process terabytes of data in 60 seconds?

You are correct that the maximum duration for a Lambda function used by S3 Object Lambda is 60 seconds.

Given the time constraint, it is not feasible to process terabytes of data within a single Lambda function execution.

S3 Object Lambda is designed for lightweight and real-time transformations rather than extensive processing of large datasets.

To handle terabytes of data, you would typically need to implement a distributed processing solution using services like Amazon EMR, AWS Glue, or AWS Batch. These services are specifically designed to handle big data workloads and provide scalability and distributed processing capabilities.

So, while S3 Object Lambda can be useful for lightweight processing tasks, it is not the appropriate tool for processing terabytes of data within the execution time limits of a Lambda function.

upvoted 2 times

Kp88 1 year, 10 months ago

Terabyte is just the storage. Lambda only need to process which application request. Think like removing/scratching off your social security number before sharing your doc to a third party.

upvoted 3 times

kruasan 2 years, 1 month ago

Selected Answer: B

- Storing the raw data in S3 provides a durable, scalable data lake. S3 requires little ongoing management overhead.
- S3 Object Lambda can be used to filter and process the data on retrieval transparently. This minimizes operational overhead by avoiding the need to preprocess and store multiple transformed copies of the data.
- Only one copy of the data needs to be stored and maintained in S3. S3 Object Lambda will transform the data on read based on the requesting application.
- No additional applications or proxies need to be developed and managed to handle the data transformation. S3 Object Lambda provides this functionality.

upvoted 4 times

kruasan 2 years, 1 month ago

Option A requires developing and managing a proxy app layer to handle data transformation, adding overhead. Options C and D require preprocessing and storing multiple copies of the transformed data, adding storage and management overhead.

Option B using S3 Object Lambda minimizes operational overhead by handling data transformation on read transparently using the native S3 functionality. Only one raw data copy is stored in S3, with no additional applications required.

upvoted 2 times

pagom 2 years, 3 months ago

Selected Answer: B

<https://aws.amazon.com/ko/blogs/korea/introducing-amazon-s3-object-lambda-use-your-code-to-process-data-as-it-is-being-retrieved-from-s3/>

upvoted 5 times

LuckyAro 2 years, 3 months ago

Selected Answer: B

B is the correct answer.

Amazon S3 Object Lambda allows you to add custom code to S3 GET requests, which means that you can modify the data before it is returned to the requesting application. In this case, you can use S3 Object Lambda to remove the PII before the data is returned to the two applications that do not need to process PII. This approach has the least operational overhead because it does not require creating separate datasets or proxy application layers, and it allows you to maintain a single copy of the data in an S3 bucket.

upvoted 5 times

NolaHOLA 2 years, 3 months ago

To meet the requirement of removing the PII before processing by two of the applications, it would be most efficient to use

To meet the requirement of removing the PII before processing by two of the applications, it would be most efficient to use option B, which involves storing the data in an Amazon S3 bucket and using S3 Object Lambda to process and transform the data before returning it to the requesting application. This approach allows the PII to be removed in real-time and without the need to create separate datasets or tables for each application. S3 Object Lambda can be configured to automatically remove PII from the data before it is sent to the non-PII processing applications. This solution provides a cost-effective and scalable way to meet the requirement with the least operational overhead.

upvoted 4 times

minglu 2 years, 3 months ago

Selected Answer: B

I think it is B.

upvoted 2 times

skiwili 2 years, 3 months ago

Selected Answer: C

Looks like C is the correct answer

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #296

Topic 1

A development team has launched a new application that is hosted on Amazon EC2 instances inside a development VPC. A solutions architect needs to create a new VPC in the same account. The new VPC will be peered with the development VPC. The VPC CIDR block for the development VPC is 192.168.0.0/24. The solutions architect needs to create a CIDR block for the new VPC. The CIDR block must be valid for a VPC peering connection to the development VPC.


What is the SMALLEST CIDR block that meets these requirements?

- A. 10.0.1.0/32
- B. 192.168.0.0/24
- C. 192.168.1.0/32

D. 10.0.1.0/24 **Most Voted**

Correct Answer: D

Community vote distribution



Comments

BrainOBrian **Highly Voted** 1 year, 9 months ago

Selected Answer: D

10.0.1.0/32 and 192.168.1.0/32 are too small for VPC, and /32 network is only 1 host
192.168.0.0/24 is overlapping with existing VPC

upvoted 30 times

kruasan **Highly Voted** 1 year, 7 months ago

Selected Answer: D

- Option A (10.0.1.0/32) is invalid - a /32 CIDR prefix is a host route, not a VPC range.
- Option B (192.168.0.0/24) overlaps the development VPC and so cannot be used.
- Option C (192.168.1.0/32) is invalid - a /32 CIDR prefix is a host route, not a VPC range.
- Option D (10.0.1.0/24) satisfies the non-overlapping CIDR requirement but is a larger block than needed. Since only two VPCs need to be peered, a /24 block provides more addresses than necessary.

upvoted 13 times

TheFivePips **Most Recent** 9 months, 2 weeks ago

Selected Answer: D

In an Amazon VPC, the first four and the last IP address in each subnet are reserved for specific purposes, and they cannot be used for customer instances. Here's how the reserved addresses are typically allocated:

Network Address (First IP):

The first IP address (all zeros in the host portion) in a subnet is reserved as the network address. For example, if you have a subnet with a CIDR notation of 10.0.0.0/24, the network address would be 10.0.0.0.

VPC Router (Second IP):

The second IP address in the subnet is reserved for the VPC router.

DNS Server (Third IP):

The third IP address is reserved for the DNS server.

Reserved for Future Use (Fourth IP):

The fourth IP address is reserved for future use.

Customer Instances (Fifth to Second-to-Last IP):

The IP addresses from the fifth to the second-to-last IP address in the subnet are available for customer instances.

Broadcast Address (Last IP):

The last IP address (all ones in the host portion) in a subnet is reserved as the broadcast address, even though AWS does not support broadcast.

upvoted 6 times

walter9660 9 months, 2 weeks ago

Selected Answer: C

10.0.0.0 - 10.255.255.255 (10/8 prefix): Example CIDR block: 10.0.0.0/16

172.16.0.0 - 172.31.255.255 (172.16/12 prefix): Example CIDR block: 172.31.0.0/16

192.168.0.0 - 192.168.255.255 (192.168/16 prefix): Example CIDR block: 192.168.0.0/20

Given that the development VPC already uses 192.168.0.0/24, we need to choose a non-overlapping CIDR block. The smallest valid CIDR block that meets the requirements is 192.168.1.0/24 (Option C).

upvoted 1 times

bora4motion 1 month, 1 week ago

WOW! just WOW!

upvoted 1 times

Murtadhaceit 1 year ago

Selected Answer: D

A and C are host IP addresses.

B is not possible because it's using the same subnet for the other team/department.

We are left with D, which is the right answer.

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

10.0.1.0/32 and 192.168.1.0/32 are too small for VPC, and /32 network is only 1 host

192.168.0.0/24 is overlapping with existing VPC

upvoted 2 times

Abrar2022 1 year, 6 months ago

Definitely D. The only valid VPC CIDR block that does not overlap with the development VPC CIDR block among the options. The other 2 CIDR block options are too small.

upvoted 2 times

antropaws 1 year, 6 months ago

Selected Answer: D

D is correct.

upvoted 2 times

channn 1 year, 8 months ago

Selected Answer: D

D is the only correct answer

upvoted 2 times

r04dB10ck 1 year, 8 months ago

Selected Answer: D

only one valid with no overlap

upvoted 2 times

Steve_4542636 1 year, 9 months ago

Selected Answer: D

A process by elimination solution here. a CIDR value is the number of bits that are locked so 10.0.0.0/32 means no range.

upvoted 4 times

LuckyAro 1 year, 9 months ago

Selected Answer: D

Answer is D, 10.0.1.0/24.

upvoted 2 times

skiwili 1 year, 9 months ago

Selected Answer: D

Yes D is the answer

upvoted 2 times

obatunde 1 year, 9 months ago

Selected Answer: D

Definitely D. It is the only valid VPC CIDR block that does not overlap with the development VPC CIDR block among the options.

upvoted 2 times

bdp123 1 year, 9 months ago

Selected Answer: D

The allowed block size is between a /28 netmask and /16 netmask.

The CIDR block must not overlap with any existing CIDR block that's associated with the VPC.

<https://docs.aws.amazon.com/vpc/latest/userguide/configure-your-vpc.html>

upvoted 6 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #297

Topic 1

A company deploys an application on five Amazon EC2 instances. An Application Load Balancer (ALB) distributes traffic to the instances by using a target group. The average CPU usage on each of the instances is below 10% most of the time, with occasional surges to 65%.

A solutions architect needs to implement a solution to automate the scalability of the application. The solution must optimize the cost of the architecture and must ensure that the application has enough CPU resources when surges occur.

Which solution will meet these requirements?

A. Create an Amazon CloudWatch alarm that enters the ALARM state when the CPUUtilization metric is less than 20%. Create an AWS Lambda function that the CloudWatch alarm invokes to terminate one of the EC2 instances in the ALB target group.


B. Create an EC2 Auto Scaling group. Select the existing ALB as the load balancer and the existing target group as the target group. Set a target tracking scaling policy that is based on the ASGAverageCPUUtilization metric. Set the minimum instances to 2, the desired capacity to 3, the maximum instances to 6, and the target value to 50%. Add the EC2 instances to the Auto Scaling group. **Most Voted**

C. Create an EC2 Auto Scaling group. Select the existing ALB as the load balancer and the existing target group as the target group. Set the minimum instances to 2, the desired capacity to 3, and the maximum instances to 6. Add the EC2 instances to the Auto Scaling group.

D. Create two Amazon CloudWatch alarms. Configure the first CloudWatch alarm to enter the ALARM state when the average CPUUtilization metric is below 20%. Configure the second CloudWatch alarm to enter the ALARM state when the average CPUUtilization metric is above 50%. Configure the alarms to publish to an Amazon Simple Notification Service (Amazon SNS) topic to send an email message. After receiving the message, log in to decrease or increase the number of EC2 instances that are running.

Correct Answer: B

Community vote distribution



Comments

bdp123 **Highly Voted** 1 year, 9 months ago

Selected Answer: B

Just create an auto scaling policy

upvoted 17 times

vilagiri **Most Recent** 1 year, 2 months ago

I picked B.. I am not 100% sure..The application is deployed in 5 instances initially. What is the logic behind 2/3/6 ASG. Because utilization is 10%, we can set min 2? I know for sure I am not going to get this ASG question correct in the exam.

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

The correct answer is B.

This solution will meet the requirements because it will:

Automate the scalability of the application by using EC2 Auto Scaling.

Optimize the cost of the architecture by only scaling the number of EC2 instances up when needed.

Ensure that the application has enough CPU resources when surges occur by setting the target value of the target tracking scaling policy to 50%.

upvoted 3 times

ajchi1980 1 year, 5 months ago

Wrong answers: Options A, C, and D are not the most appropriate solutions:

Option A suggests creating a CloudWatch alarm to terminate an EC2 instance when CPU utilization is less than 20%. However, this approach does not ensure that the application will have enough CPU resources during surges, as it only terminates instances when CPU utilization is low, which may not meet the requirements.

Option C suggests creating an Auto Scaling group without any specific scaling policies or configurations. This approach does not address the need for automated scaling based on CPU utilization, making it insufficient for the given requirements.

Option D suggests using CloudWatch alarms to send notifications via Amazon SNS and manually adjusting the number of instances based on the received messages. This approach lacks automation and requires manual intervention, which does not optimize cost or meet the requirement of automated scalability.

Therefore, Option B is the most appropriate solution in this case.

upvoted 3 times

ajchi1980 1 year, 5 months ago

Selected Answer: B

Explanation:

Option B leverages EC2 Auto Scaling, which is designed to automatically adjust the number of instances based on specified metrics. By setting a target tracking scaling policy based on average CPU utilization, the Auto Scaling group can dynamically scale the number of instances to maintain the desired level of CPU resources. The minimum instances of 2 ensure a minimum baseline capacity, while the desired capacity of 3 ensures at least three instances are running even during normal traffic. The maximum instances of 6 cap the upper limit to control costs.

upvoted 4 times

RoroJ 1 year, 6 months ago

Selected Answer: D

Auto Scaling group must have an AMI for it.

upvoted 2 times

pentium75 11 months, 2 weeks ago

"After receiving the message, log in to decrease or increase the number of EC2 instances that are running" does surely not "automate the scalability".

upvoted 3 times

th3k33n 1 year, 6 months ago

how can we set max to 6 since the company is using 5 ec2 instance

upvoted 2 times

examtopictempacc 1 year, 6 months ago

In the scenario you provided, you're setting up an Auto Scaling group to manage the instances for you, and the settings (min 2, desired 3, max 6) are for the Auto Scaling group, not for your existing instances. When you integrate the instances into the Auto Scaling group, you are effectively moving from a fixed instance count to a dynamic one that can range from 2 to 6 based on the demand.

The existing 5 instances can be included in the Auto Scaling group, but the group can reduce the number of instances if the load is low (to the minimum specified, which is 2 in this case) and can also add more instances (up to a maximum of 6) if the load increases.

upvoted 4 times

kruasan 1 year, 7 months ago

Selected Answer: B

Reasons:

- An Auto Scaling group will automatically scale the EC2 instances to match changes in demand. This optimizes cost by only running as many instances as needed.
- A target tracking scaling policy monitors the ASGAverageCPUUtilization metric and scales to keep the average CPU around the 50% target value. This ensures there are enough resources during CPU surges.
- The ALB and target group are reused, so the application architecture does not change. The Auto Scaling group is associated to the existing load balancer setup.
- A minimum of 2 and maximum of 6 instances provides the ability to scale between 3 and 6 instances as needed based on demand.
- Costs are optimized by starting with only 3 instances (the desired capacity) and scaling up as needed. When CPU usage drops, instances are terminated to match the desired capacity.

upvoted 3 times

kruasan 1 year, 7 months ago

Option A - terminates instances reactively based on low CPU and may not provide enough capacity during surges. Does not optimize cost.

Option C - lacks a scaling policy so will not automatically adjust capacity based on changes in demand. Does not ensure enough resources during surges.

Option D - requires manual intervention to scale capacity. Does not optimize cost or provide an automated solution.

upvoted 2 times

darn 1 year, 7 months ago

as you dig down the question, they get more and more bogus with less and less votes

upvoted 1 times

Steve_4542636 1 year, 9 months ago

Selected Answer: B

B is my vote

upvoted 2 times

KZM 1 year, 9 months ago

Based on the information given, the best solution is option "B".

Autoscaling group with target tracking scaling policy with min 2 instances, desired capacity to 3, and the maximum instances to 6.

upvoted 2 times

Shrestwt 1 year, 7 months ago

But the company is using only 5 EC2 Instances so how can we set maximum instance to 6.

upvoted 2 times

pentium75 11 months, 2 weeks ago

"Create an EC2 Auto Scaling group" includes replacing your existing EC2 instances with a launch configuration that starts and stops instances automatically.

upvoted 2 times

LuckyAro 1 year, 9 months ago

Selected Answer: B

B is the correct solution because it allows for automatic scaling based on the average CPU utilization of the EC2 instances in the target group. With the use of a target tracking scaling policy based on the ASGAverageCPUUtilization metric, the EC2 Auto Scaling group can ensure that the target value of 50% is maintained while scaling the number of instances in the group up or

down as needed. This will help ensure that the application has enough CPU resources during surges without overprovisioning, thus optimizing the cost of the architecture.

upvoted 3 times

Babba 1 year, 9 months ago

Selected Answer: B

Should be B

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #298

Topic 1

A company is running a critical business application on Amazon EC2 instances behind an Application Load Balancer. The EC2 instances run in an Auto Scaling group and access an Amazon RDS DB instance.

The design did not pass an operational review because the EC2 instances and the DB instance are all located in a single Availability Zone. A solutions architect must update the design to use a second Availability Zone.

Which solution will make the application highly available?

- A. Provision a subnet in each Availability Zone. Configure the Auto Scaling group to distribute the EC2 instances across both Availability Zones. Configure the DB instance with connections to each network.
- B. Provision two subnets that extend across both Availability Zones. Configure the Auto Scaling group to distribute the EC2 instances across both Availability Zones. Configure the DB instance with connections to each network.
- C. Provision a subnet in each Availability Zone. Configure the Auto Scaling group to distribute the EC2 instances across both Availability Zones. Configure the DB instance for Multi-AZ deployment. **Most Voted**
- D. Provision a subnet that extends across both Availability Zones. Configure the Auto Scaling group to distribute the EC2 instances across both Availability Zones. Configure the DB instance for Multi-AZ deployment.

Correct Answer: C

Community vote distribution

C (100%)

Comments

bdp123 **Highly Voted** 1 year, 9 months ago

Selected Answer: C

A subnet must reside within a single Availability Zone.

<https://aws.amazon.com/vpc/faqs/#:~:text=Can%20a%20subnet%20span%20Availability,within%20a%20single%20Availability%20Zone.>

upvoted 23 times

zjcorpuz **Highly Voted** 1 year, 4 months ago

a subnet only resides on a one AZ, it does not span to another AZ.

upvoted 7 times

upvoted 7 times

LoXoL Most Recent 10 months, 1 week ago

Selected Answer: C

A subnet can't "Extend" across multiple AZs: B,D out
HA = RDS Multi-AZ: A out

C

upvoted 4 times

thewalker 10 months, 1 week ago

Selected Answer: C

An Auto Scaling group can span across two Availability Zones, where one subnet is created in each AZ.

When creating an Auto Scaling group, you need to specify at least one subnet. You can add additional subnets later on, including subnets across multiple AZs.

Auto Scaling will distribute instances evenly across the specified subnets to maintain availability and optimize performance. If one AZ becomes unavailable, instances can be launched in the other AZ.

The associated load balancer should also span the same subnets/AZs as the Auto Scaling group. This allows traffic to be routed to instances in different subnets and AZs, increasing fault tolerance of the application.

upvoted 3 times

TariqKipkemei 1 year, 2 months ago

Selected Answer: C

Provision a subnet in each Availability Zone. Configure the Auto Scaling group to distribute the EC2 instances across both Availability Zones. Configure the DB instance for Multi-AZ deployment

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

This solution will ensure that the EC2 instances and the DB instance are not located in the same Availability Zone, which will improve the availability of the application.

upvoted 3 times

MrAWSAssociate 1 year, 5 months ago

Selected Answer: C

D is completely wrong, because each subnet must reside entirely within one Availability Zone and cannot span zones. By launching AWS resources in separate Availability Zones, you can protect your applications from the failure of a single Availability Zone.

upvoted 3 times

Anmol_1010 1 year, 5 months ago

The key word here was extend.

upvoted 2 times

GalileoEC2 1 year, 8 months ago

This discards B and D: Subnet basics. Each subnet must reside entirely within one Availability Zone and cannot span zones. By launching AWS resources in separate Availability Zones, you can protect your applications from the failure of a single Availability Zone

upvoted 3 times

Steve_4542636 1 year, 9 months ago

Selected Answer: C

a subnet is per AZ. a scaling group can span multiple AZs. <https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-add-availability-zone.html>

upvoted 2 times

KZM 1 year, 9 months ago

I think D.

Span the single subnet in both Availability Zones can access the DB instances in either zone without going over the public internet.
upvoted 3 times

KZM 1 year, 9 months ago

Can span like that?
upvoted 1 times

KZM 1 year, 9 months ago

Sorry I think C is correct.
upvoted 2 times

leoattf 1 year, 9 months ago

Nope. The answer is indeed C.
You cannot span like that. Check the link below:
"Each subnet must reside entirely within one Availability Zone and cannot span zones."
<https://docs.aws.amazon.com/vpc/latest/userguide/configure-subnets.html>
upvoted 4 times

KZM 1 year, 9 months ago

Thanks, Leoattf for the link you shared.
upvoted 3 times

Babba 1 year, 9 months ago

Selected Answer: C

it's C
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #299

Topic 1

A research laboratory needs to process approximately 8 TB of data. The laboratory requires sub-millisecond latencies and a minimum throughput of 6 GBps for the storage subsystem. Hundreds of Amazon EC2 instances that run Amazon Linux will distribute and process the data.

Which solution will meet the performance requirements?

- A. Create an Amazon FSx for NetApp ONTAP file system. Set each volume's tiering policy to ALL. Import the raw data into the file system. Mount the file system on the EC2 instances.
- B. Create an Amazon S3 bucket to store the raw data. Create an Amazon FSx for Lustre file system that uses persistent SSD storage. Select the option to import data from and export data to Amazon S3. Mount the file system on the EC2 instances. **Most Voted**
- C. Create an Amazon S3 bucket to store the raw data. Create an Amazon FSx for Lustre file system that uses persistent HDD storage. Select the option to import data from and export data to Amazon S3. Mount the file system on the EC2 instances.
- D. Create an Amazon FSx for NetApp ONTAP file system. Set each volume's tiering policy to NONE. Import the raw data into the file system. Mount the file system on the EC2 instances.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**Bhawesh** **Highly Voted** 1 year, 9 months ago**Selected Answer: B**

Keyword here is a minimum throughput of 6 GBps. Only the FSx for Lustre with SSD option gives the sub-milli response and throughput of 6 GBps or more.

B. Create an Amazon S3 bucket to store the raw data. Create an Amazon FSx for Lustre file system that uses persistent SSD storage. Select the option to import data from and export data to Amazon S3. Mount the file system on the EC2 instances.

References:

<https://aws.amazon.com/fsx/when-to-choose-fsx/>

upvoted 20 times

bdp123 Highly Voted 1 year, 9 months ago

Selected Answer: B

Create an Amazon S3 bucket to store the raw data Create an Amazon FSx for Lustre file system that uses persistent SSD storage Select the option to import data from and export data to Amazon S3 Mount the file system on the EC2 instances. Amazon FSx for Lustre uses SSD storage for submillisecond latencies and up to 6 GBps throughput, and can import data from and export data to Amazon S3. Additionally, the option to select persistent SSD storage will ensure that the data is stored on the disk and not lost if the file system is stopped.

upvoted 8 times

MrPCarrot Most Recent 9 months, 2 weeks ago

Answer is B : FSx for Lustre with SSD option gives the sub-milli response and throughput of 6 GBps or more

upvoted 3 times

Pangian 9 months, 4 weeks ago

I dont even think that NetApp comes for Linux

upvoted 1 times

djgodzilla 11 months, 1 week ago

Selected Answer: B

Amazon FSx for Lustre for compute-intensive workloads.

- allows file-based applications to access data with hundreds of gigabytes per second of data, millions of IOPS, and sub millisecond latencies.
- supports file access to thousands of EC2 instances

and well SSD always wins ;)

upvoted 3 times

Mikado211 12 months ago

Selected Answer: B

sub-millisecondes == Lustre

HDD vs SSD == for performance use SSD

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Amazon FSx for Lustre with SSD: Amazon FSx for Lustre is designed for high-performance, parallel file processing workloads. Choosing SSD storage ensures fast I/O and meets the sub-millisecond latency requirement.

upvoted 3 times

rolervengador 1 year, 3 months ago

Voto por la B

upvoted 2 times

Gooniegoogoo 1 year, 5 months ago

So many of these are wrong, its good we have people that vote so we can get to the right answer!!

upvoted 2 times

kruasan 1 year, 7 months ago

Selected Answer: B

- Amazon FSx for Lustre with SSD storage can provide up to 260 GB/s of aggregate throughput and sub-millisecond latencies needed for this workload.
- Persistent SSD storage ensures data durability in the file system. Data is also exported to S3 for backup storage.
- The file system will import the initial 8 TB of raw data from S3, providing a fast storage tier for processing while retaining the data in S3.
- The file system is mounted to the EC2 compute instances to distribute processing.
- FSx for Lustre is optimized for high-performance computing workloads running on Linux, matching the EC2 environment.

upvoted 3 times

kruasan 1 year, 7 months ago