

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #432

Topic 1

An ecommerce company wants to use machine learning (ML) algorithms to build and train models. The company will use the models to visualize complex scenarios and to detect trends in customer data. The architecture team wants to integrate its ML models with a reporting platform to analyze the augmented data and use the data directly in its business intelligence dashboards.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use AWS Glue to create an ML transform to build and train models. Use Amazon OpenSearch Service to visualize the data.
- B. Use Amazon SageMaker to build and train models. Use Amazon QuickSight to visualize the data. **Most Voted**
- C. Use a pre-built ML Amazon Machine Image (AMI) from the AWS Marketplace to build and train models. Use Amazon OpenSearch Service to visualize the data.
- D. Use Amazon QuickSight to build and train models by using calculated fields. Use Amazon QuickSight to visualize the data.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**vincent2023** 4 months, 2 weeks ago**Selected Answer: B**

Sage Maker - ML
upvoted 1 times

67db0ed 10 months, 1 week ago

<https://docs.aws.amazon.com/quicksight/latest/user/sagemaker-integration.html>
upvoted 2 times

awsgeek75 1 year, 5 months ago**Selected Answer: B**

Machine Learning = Sage Maker so B for least operational overhead

A and D are not right technologies.

C is possible but with more overhead of using AMI even if you can get OpenSearch to visualize the data somehow which I don't think is possible without massive overhead

upvoted 4 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

Use Amazon SageMaker to build and train models. Use Amazon QuickSight to visualize the data.

upvoted 3 times

james2033 1 year, 10 months ago

Selected Answer: B

Question keyword "machine learning", answer keyword "Amazon SageMaker". Choose B. Use Amazon QuickSight for visualization. See "Gaining insights with machine learning (ML) in Amazon QuickSight" at

<https://docs.aws.amazon.com/quicksight/latest/user/making-data-driven-decisions-with-ml-in-quicksight.html>

upvoted 3 times

VellaDevil 1 year, 11 months ago

Selected Answer: B

Sagemaker.

upvoted 2 times

TariqKipkemei 2 years ago

Selected Answer: B

Business intelligence, visualiations = AmazonQuicksight

ML = Amazon SageMaker

upvoted 3 times

antropaws 2 years ago

Selected Answer: B

Most likely B.

upvoted 2 times

omoakin 2 years ago

Amazon SageMaker is a fully managed service that provides every developer and data scientist with the ability to build, train, and deploy ML models quickly.

upvoted 2 times

cloudenthusiast 2 years ago

Amazon SageMaker is a fully managed service that provides a complete set of tools and capabilities for building, training, and deploying ML models. It simplifies the end-to-end ML workflow and reduces operational overhead by handling infrastructure provisioning, model training, and deployment.

To visualize the data and integrate it into business intelligence dashboards, Amazon QuickSight can be used. QuickSight is a cloud-native business intelligence service that allows users to easily create interactive visualizations, reports, and dashboards from various data sources, including the augmented data generated by the ML models.

upvoted 4 times

Efren 2 years ago

Selected Answer: B

ML== SageMaker

upvoted 2 times

nosense 2 years ago

Selected Answer: B

B sagemaker provide deploy ml models

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #433

Topic 1

A company is running its production and nonproduction environment workloads in multiple AWS accounts. The accounts are in an organization in AWS Organizations. The company needs to design a solution that will prevent the modification of cost usage tags.

Which solution will meet these requirements?

- A. Create a custom AWS Config rule to prevent tag modification except by authorized principals.
- B. Create a custom trail in AWS CloudTrail to prevent tag modification.
- C. Create a service control policy (SCP) to prevent tag modification except by authorized principals. **Most Voted**
- D. Create custom Amazon CloudWatch logs to prevent tag modification.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Guru4Cloud** **Highly Voted** 1 year, 3 months ago**Selected Answer: C**

Tip: AWS Organizaton + service control policy (SCP) - This for any questions, you see both together. then you tell me C. Create a service control policy (SCP) to prevent tag modification except by authorized principals.

upvoted 7 times

awsgeek75 **Most Recent** 11 months ago**Selected Answer: C**

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps.html

AWS example for this question/use case:

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps_examples_tagging.html#example-require-restrict-tag-mods-to-admin

upvoted 3 times

james2033 1 year, 4 months ago**Selected Answer: C**

D "Amazon CloudWatch" just for logging, not for prevent tag modification
https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_tag-policies-cwe.html

Amazon Organziaton has "Service Control Policy (SCP)" with "tag policy"
https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_tag-policies.html . Choose C.

AWS Config for technical stuff, not for tag policies. Not A.
upvoted 4 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: C

Service control policies (SCPs) are a type of organization policy that you can use to manage permissions in your organization.
upvoted 2 times

alexandercamachop 1 year, 6 months ago

Selected Answer: C

Anytime we need to restrict anything in an AWS Organization, it is SCP Policies.
upvoted 3 times

Abrar2022 1 year, 6 months ago

AWS Config is for tracking configuration changes
upvoted 2 times

Abrar2022 1 year, 6 months ago

so it's wrong. Right asnwer is C
upvoted 3 times

antropaws 1 year, 6 months ago

Selected Answer: C

I'd say C.
upvoted 3 times

hiroohiroy 1 year, 6 months ago

Selected Answer: C

https://docs.aws.amazon.com/ja_jp/organizations/latest/userguide/orgs_manage_policies_scps_examples_tagging.html
upvoted 4 times

nosense 1 year, 6 months ago

Selected Answer: C

Denies tag: modify
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #434

Topic 1

A company hosts its application in the AWS Cloud. The application runs on Amazon EC2 instances behind an Elastic Load Balancer in an Auto Scaling group and with an Amazon DynamoDB table. The company wants to ensure the application can be made available in another AWS Region with minimal downtime.

What should a solutions architect do to meet these requirements with the LEAST amount of downtime?

- A. Create an Auto Scaling group and a load balancer in the disaster recovery Region. Configure the DynamoDB table as a global table. Configure DNS failover to point to the new disaster recovery Region's load balancer. **Most Voted**
- B. Create an AWS CloudFormation template to create EC2 instances, load balancers, and DynamoDB tables to be launched when needed. Configure DNS failover to point to the new disaster recovery Region's load balancer.
- C. Create an AWS CloudFormation template to create EC2 instances and a load balancer to be launched when needed. Configure the DynamoDB table as a global table. Configure DNS failover to point to the new disaster recovery Region's load balancer.
- D. Create an Auto Scaling group and load balancer in the disaster recovery Region. Configure the DynamoDB table as a global table. Create an Amazon CloudWatch alarm to trigger an AWS Lambda function that updates Amazon Route 53 pointing to the disaster recovery load balancer.

Correct Answer: A*Community vote distribution***Comments****lucdt4** **Highly Voted** 2 years ago**Selected Answer: A**

A and D is correct.

But Route 53 has a feature DNS failover when instances down so we don't need use Cloudwatch and lambda to trigger
-> A correct

upvoted 16 times

Wablo 1 year, 11 months ago

Yes it does but you configure it. It's not automated anymore. D is the best answer!

upvoted 2 times

JA2018 6 months, 2 weeks ago

to be launched when needed => requires manual human intervention

Certainly will not help with requirement for the least amount of downtime.

upvoted 1 times

Kp88 1 year, 10 months ago

What are you talking about configuring ? Yes you have to configure everything at some point
<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/dns-failover-configuring.html>

upvoted 2 times

smartegnine 1 year, 11 months ago

Did not see Route 53 in this question right? So my opinion is D

upvoted 1 times

pentium75 Highly Voted 1 year, 5 months ago

Selected Answer: C

They are not asking for automatic failover, they want to "ensure the application can (!) be made available in another AWS Region with minimal downtime". This works with C; they would just execute the template and it would be available in short time.

A would create a DR environment that IS already available, which is not what the question asks for.

D is like A, just abusing Lambda to update the DNS record (which doesn't make sense).

B would create a separate, empty database

upvoted 9 times

bora4motion 1 month, 1 week ago

I am going with A. There's no downtime with A, maybe a blip because of DNS.

upvoted 1 times

[Removed] Most Recent 10 months, 1 week ago

ChatGPT:

Option C involves creating an AWS CloudFormation template to create EC2 instances and a load balancer only when needed, and configuring the DynamoDB table as a global table. This approach might introduce more downtime because the infrastructure in the disaster recovery region is not pre-deployed and ready to take over immediately. The process of launching instances and configuring the load balancer can take some time, leading to delays during the failover.

Option A, on the other hand, ensures that the necessary infrastructure (Auto Scaling group, load balancer, and DynamoDB global table) is already set up and running in the disaster recovery region. This pre-deployment reduces downtime since the failover can be handled quickly by updating DNS to point to the disaster recovery region's load balancer.

upvoted 2 times

anikolov 1 year, 4 months ago

Selected Answer: A

With the LEAST amount of downtime = A

Cost effective = C, but risky some of EC2 types/capacity not available in Region at the time, when need to switch to DR

upvoted 7 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

There are 2 parts. DB and application. Dynamo DB recovery in another region is not possible without global table so option B is out.

A will make the infra available in 2 regions which is not required. The question is about DR, not scaling.

D Use Lambda to modify R53 to point to new region. This is going to cause delays but is possible and it will also be running a scaled EC2 instances in passive region.

C Make a CF template which can launch the infra when needed. DB is global table so it will be available.

upvoted 4 times

bora4motion 1 month, 1 week ago

you need minimum downtime, nobody asks about what's required or not apart of minimal downtime. I'm going with a.
upvoted 1 times

LeonSauveterre 6 months, 1 week ago

Option C relies on launching EC2 instances and load balancers only when needed, which may increase downtime. On contrary, I think A is better because it ensures resources are pre-configured and ready-to-go.

It's true nobody mentioned anything about infra availability in 2 regions... but it couldn't hurt since we simply need the least amount of downtime.
upvoted 1 times

meowruki 1 year, 6 months ago

Selected Answer: C

AWS CloudFormation Template: Use CloudFormation to define the infrastructure components (EC2 instances, load balancer, etc.) in a template. This allows for consistent and repeatable infrastructure deployment.

EC2 Instances and Load Balancer: Launch the EC2 instances and load balancer in the disaster recovery (DR) Region using the CloudFormation template. This enables the deployment of the application in the DR Region when needed.

DynamoDB Global Table: Configure the DynamoDB table as a global table. DynamoDB Global Tables provide automatic multi-region, multi-master replication, ensuring that the data is available in both the primary and DR Regions.

DNS Failover: Configure DNS failover to point to the new DR Region's load balancer. This allows for seamless failover of traffic to the DR Region when needed.

Option A is close, but it introduces an Auto Scaling group in the disaster recovery Region, which might introduce unnecessary complexity and potential scaling delays. Option D introduces a Lambda function triggered by CloudWatch alarms, which might add latency and complexity compared to the more direct approach in Option C.
upvoted 1 times

bogobob 1 year, 6 months ago

Selected Answer: A

Assuming theyre using Route53 as a DNS then A <https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/dns-failover.html>
upvoted 3 times

EEK2k 1 year, 6 months ago

Selected Answer: C

Only B and C take care of EC2 instants. But since B does not take care of Data in the Dynamo DB, C is the only correct Answer.
upvoted 1 times

potomac 1 year, 7 months ago

Selected Answer: A

Route 53 haves a feature DNS failover when instances down
upvoted 1 times

thanhnv142 1 year, 7 months ago

C is the best choice here
upvoted 1 times

Wayne23Fang 1 year, 7 months ago

Selected Answer: C

I think CloudFormation is easier than manual provision of Auto Scaling group and load balancer in DR region.
upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

Creating Auto Scaling group and load balancer in DR region allows fast launch of capacity when needed.
Configuring DynamoDB as a global table provides continuous data replication

Configuring DynamoDB as a global table provides continuous data replication.
Using DNS failover via Route 53 to point to the DR region's load balancer enables rapid traffic shifting.

upvoted 2 times

Wablo 1 year, 11 months ago

Both Option A and Option D include the necessary steps of setting up an Auto Scaling group and load balancer in the disaster recovery Region, configuring the DynamoDB table as a global table, and updating DNS records. However, Option D provides a more detailed approach by explicitly mentioning the use of an Amazon CloudWatch alarm and AWS Lambda function to automate the DNS update process.

By leveraging an Amazon CloudWatch alarm, Option D allows for an automated failover mechanism. When triggered, the CloudWatch alarm can execute an AWS Lambda function, which in turn can update the DNS records in Amazon Route 53 to redirect traffic to the disaster recovery load balancer in the new Region. This automation helps reduce the potential for human error and further minimizes downtime.

Answer is D

upvoted 2 times

Kp88 1 year, 10 months ago

Failover policy takes care of DNS record update so no need for cloud watch/lambda

upvoted 1 times

TariqKipkemei 1 year, 12 months ago

Selected Answer: C

The company wants to ensure the application 'CAN' be made available in another AWS Region with minimal downtime. Meaning they want to be able to launch infra on need basis.

Best answer is C.

upvoted 2 times

Wablo 1 year, 11 months ago

minimal downtime not minimal effort!

D

upvoted 1 times

dajform 1 year, 11 months ago

B, C are not OK because "launching resources when needed", which will increase the time to recover "DR"

upvoted 2 times

AshishRocks 2 years ago

I feel it is A

Configure DNS failover: Use DNS failover to point the application's DNS record to the load balancer in the disaster recovery Region. DNS failover allows you to route traffic to the disaster recovery Region in case of a failure in the primary Region.

upvoted 3 times

Wablo 1 year, 11 months ago

Once you configure manually the DNS, it's no more automated like Lambda does.

upvoted 2 times

Yadav_Sanjay 2 years ago

Selected Answer: C

C suits best

upvoted 3 times

hiroohiroo 2 years ago

Selected Answer: A

AがDNS フェイルオーバー

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #435

Topic 1


A company needs to migrate a MySQL database from its on-premises data center to AWS within 2 weeks. The database is 20 TB in size. The company wants to complete the migration with minimal downtime.

Which solution will migrate the database MOST cost-effectively?

- A. Order an AWS Snowball Edge Storage Optimized device. Use AWS Database Migration Service (AWS DMS) with AWS Schema Conversion Tool (AWS SCT) to migrate the database with replication of ongoing changes. Send the Snowball Edge device to AWS to finish the migration and continue the ongoing replication. **Most Voted**
- B. Order an AWS Snowmobile vehicle. Use AWS Database Migration Service (AWS DMS) with AWS Schema Conversion Tool (AWS SCT) to migrate the database with ongoing changes. Send the Snowmobile vehicle back to AWS to finish the migration and continue the ongoing replication.
- C. Order an AWS Snowball Edge Compute Optimized with GPU device. Use AWS Database Migration Service (AWS DMS) with AWS Schema Conversion Tool (AWS SCT) to migrate the database with ongoing changes. Send the Snowball device to AWS to finish the migration and continue the ongoing replication
- D. Order a 1 GB dedicated AWS Direct Connect connection to establish a connection with the data center. Use AWS Database Migration Service (AWS DMS) with AWS Schema Conversion Tool (AWS SCT) to migrate the database with replication of ongoing changes.

Correct Answer: A

Community vote distribution



Comments

nosense **Highly Voted** 1 year, 6 months ago

Selected Answer: A

A) 300 first 10 days. 150 shipping

D) 750 for 2 weeks

upvoted 12 times

Efren 1 year, 6 months ago

Thanks, i was checking the speed more than price. Thanks for the clarification

upvoted 3 times

Goutham4981 Highly Voted 1 year ago

Selected Answer: A

Direct Connect takes at least 1 month to setup - D is invalid

AWS Snowmobile is used for transferring large amounts of data (petabytes) from remote locations where establishing a connection to the cloud is impossible - B is invalid

AWS Snowball Edge Compute Optimized provides higher vCPU performance and lower storage as compared to Snowball storage optimized. As our need is solely data transfer, high vCPU performance is not required but high storage is - C is invalid

upvoted 7 times

zinabu Most Recent 7 months, 2 weeks ago

Selected Answer: A

But I didn't understand why we are using a schema conversion tool, because AWS have already a managed service engine for MySQL Db, (RDS for MySQL or Aurora for my SQL is on the table)

upvoted 6 times

EEK2k 1 year ago

Selected Answer: D

To calculate the time it would take to transfer 20TB of data over a 1 GB dedicated AWS Direct Connect, we can use the formula:

time = data size / data transfer rate

Here, the data size is 20TB, which is equivalent to 20,000 GB or 20,000,000 MB. The data transfer rate is 1 GB/s.

Converting the data size to MB, we get:

$20,000,000 \text{ MB} / 1 \text{ GB/s} = 20,000 \text{ seconds}$

Therefore, it would take approximately 20,000 seconds or 5.56 hours to transfer 20TB of data over a 1 GB dedicated AWS Direct Connect.

upvoted 3 times

Murtadhaceit 1 year ago

It takes more way more than 2 weeks to setup Direct Connect. Therefore, D is not valid since we have to do the transfer within 2 weeks.

upvoted 3 times

potomac 1 year, 1 month ago

Selected Answer: A

C is wrong, GPU is not needed

upvoted 4 times

Ramdi1 1 year, 2 months ago

Selected Answer: A

Has to be A. the option for D would only work if they said they have like 6 Months plus. It would take too long to set up.

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

I agreed with A.

Why not D.?

When you initiate the process by requesting an AWS Direct Connect connection, it typically starts with the AWS Direct Connect provider. This provider may need to coordinate with AWS to allocate the necessary resources. This initial setup phase can take anywhere from a few days to a couple of weeks.

Couple of weeks? No Good

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

When you create a Snowball job in the AWS console, it will estimate the delivery date based on your location. Being near a

facility shows 1-2 day estimated delivery.

For extremely urgent requests, you can contact AWS Support and inquire about expedited Snowball delivery. If inventory is available, they may be able to ship same day or next day.

upvoted 3 times

james2033 1 year, 4 months ago

Selected Answer: A

Keyword "20 TB", choose "AWS Snowball", there are A or C. C has word "GPU" what is not related, therefore choose A.

upvoted 3 times

Zox42 1 year, 5 months ago

Selected Answer: A

Answer A

upvoted 2 times

MrAWSAssociate 1 year, 5 months ago

Selected Answer: D

D is correct

upvoted 1 times

pentium75 11 months, 1 week ago

No, takes months, not weeks

upvoted 2 times

DrWatson 1 year, 6 months ago

Selected Answer: A

https://docs.aws.amazon.com/dms/latest/userguide/CHAP_LargeDBs.Process.html

upvoted 3 times

RoroJ 1 year, 6 months ago

Selected Answer: A

D Direct Connection will need a long time to setup plus need to deal with Network and Security changes with existing environment. Ad then plus the Data trans time... No way can be done in 2 weeks.

upvoted 5 times

Joselucho38 1 year, 6 months ago

Selected Answer: D

Overall, option D combines the reliability and cost-effectiveness of AWS Direct Connect, AWS DMS, and AWS SCT to migrate the database efficiently and minimize downtime.

upvoted 2 times

Abhineet9148232 1 year, 6 months ago

Selected Answer: A

D - Direct Connect takes atleast a month to setup! Requirement is for within 2 weeks.

upvoted 5 times

Rob1L 1 year, 6 months ago

Selected Answer: D

AWS Snowball Edge Storage Optimized device is used for large-scale data transfers, but the lead time for delivery, data transfer, and return shipping would likely exceed the 2-week time frame. Also, ongoing database changes wouldn't be replicated while the device is in transit.

upvoted 1 times

Rob1L 1 year, 6 months ago

Change to A because "Most cost effective"

upvoted 3 times

Rob1L 1 year, 6 months ago

nirooniroo 1 year, 6 months ago

Selected Answer: A

https://docs.aws.amazon.com/ja_jp/snowball/latest/developer-guide/device-differences.html#device-options
Aです。

upvoted 4 times

norris81 1 year, 6 months ago

Selected Answer: A

How long does direct connect take to provision ?

upvoted 2 times

examtopictempacc 1 year, 6 months ago

At least one month and expensive.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #436

Topic 1

A company moved its on-premises PostgreSQL database to an Amazon RDS for PostgreSQL DB instance. The company successfully launched a new product. The workload on the database has increased. The company wants to accommodate the larger workload without adding infrastructure.

Which solution will meet these requirements MOST cost-effectively?

- A. Buy reserved DB instances for the total workload. Make the Amazon RDS for PostgreSQL DB instance larger. **Most Voted**
- B. Make the Amazon RDS for PostgreSQL DB instance a Multi-AZ DB instance.
- C. Buy reserved DB instances for the total workload. Add another Amazon RDS for PostgreSQL DB instance.
- D. Make the Amazon RDS for PostgreSQL DB instance an on-demand DB instance.

Correct Answer: A*Community vote distribution***Comments****elmogy** **Highly Voted** 2 years ago**Selected Answer: A**

A.
"without adding infrastructure" means scaling vertically and choosing larger instance.
"MOST cost-effectively" reserved instances
upvoted 16 times

wsdadsdqwdaw 1 year, 7 months ago

"MOST cost-effectively" doesn't mean reserved instances. Only in this case it is but not in general.
upvoted 6 times

zdi561 **Most Recent** 4 months, 3 weeks ago**Selected Answer: B**

without adding infrastructure means without adding more capacity per instance, multi-az adding read replicas solve the problem.
upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

accommodate the larger workload without adding infrastructure. = Reserved DB instance
upvoted 3 times

awsgeek75 1 year, 5 months ago

+ make the instance larger so most cost effective is to reserve a large instance suitable for workload which is A
upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: A

B - Multi-AZ is for HA, does not help 'accommodating the larger workload'
C - Adding "another instance" will not help, we can't split the workload between two instances
D - On-demand instance is a good choice for unknown workload, but here we know the workload, it's just higher than before
upvoted 3 times

Goutham4981 1 year, 6 months ago

Selected Answer: A

Cannot add more infrastructure - C is invalid
Multi AZ DB instance is for high availability and failure mitigation, does not increase performance, higher workload support - B is invalid
On demand instances are costlier than Reserved instances - D is invalid
upvoted 2 times

bogobob 1 year, 6 months ago

Selected Answer: D

Not A : "launched a new product", reserved instances are for known workloads, a new product doesn't have known workload.
Not B : "accommodate the larger workload", while Multi-AZ can help with larger workloads, they are more for higher availability.
Not C : "without adding infrastructure", adding a PostGresQL instance is new infrastructure.
upvoted 3 times

JA2018 6 months, 2 weeks ago

Please refer to following Keys in Stem:

1. The company successfully launched a new product.
 2. The workload on the database has increased.
 3. The company wants to accommodate the larger workload without adding infrastructure.
- upvoted 1 times

pentium75 1 year, 5 months ago

Question says nothing about unknown load. New product -> more total products -> load has increased.
upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

B is the best approach in this scenario overall:

Making the RDS PostgreSQL instance Multi-AZ adds a standby replica to handle larger workloads and provides high availability.
Even though it adds infrastructure, the cost is less than doubling the infrastructure with a separate DB instance.
It provides better performance, availability, and disaster recovery than a single larger instance.
upvoted 2 times

BillyBlunts 1 year, 8 months ago

Agreed the answer is B
Multi-AZ deployments are cost-effective because they leverage the standby instance without incurring additional charges. You

multi-AZ deployments are cost-effective because they leverage the standby instance without incurring additional charges. You only pay for the primary instance's regular usage costs.

upvoted 1 times

pentium75 1 year, 5 months ago

Multi-AZ is for HA, does not add performance. Meaning, will not help 'accommodating the larger workload'.

upvoted 3 times

james2033 1 year, 10 months ago

Selected Answer: A

Buy larger instance.

upvoted 2 times

james2033 1 year, 10 months ago

Selected Answer: A

Keyword "Amazon RDS for PostgreSQL instance large" . See list of size of instance at <https://aws.amazon.com/rds/instance-types/>

upvoted 2 times

examtopictempacc 2 years ago

Selected Answer: A

A.

Not C: without adding infrastructure

upvoted 3 times

EA100 2 years ago

Answer - C

Option B, making the Amazon RDS for PostgreSQL DB instance a Multi-AZ DB instance, would provide high availability and fault tolerance but may not directly address the need for increased capacity to handle the larger workload.

Therefore, the recommended solution is Option C: Buy reserved DB instances for the workload and add another Amazon RDS for PostgreSQL DB instance to accommodate the increased workload in a cost-effective manner.

upvoted 1 times

cloudenthusiast 2 years ago

C

Option C: buying reserved DB instances for the total workload and adding another Amazon RDS for PostgreSQL DB instance seems to be the most appropriate choice. It allows for workload distribution across multiple instances, providing scalability and potential performance improvements. Additionally, reserved instances can provide cost savings in the long term.

upvoted 1 times

nosense 2 years ago

A for me, because without adding additional infrastructure

upvoted 4 times

th3k33n 2 years ago

Should be C

upvoted 1 times

Efren 2 years ago

That would add more infrastructure. A would increase the size, keeping the number of instances, i think

upvoted 2 times

cloudenthusiast 2 years ago

Option A involves making the existing Amazon RDS for PostgreSQL DB instance larger. While this can improve performance, it may not be sufficient to handle a significantly increased workload. It also doesn't distribute the workload or provide scalability.

upvoted 1 times

nosense 2 years ago

husense 2 years ago

The main not HA, cost-effectively and without adding infrastructure
upvoted 2 times

omoakin 2 years ago

A is the best
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #437

Topic 1


A company operates an ecommerce website on Amazon EC2 instances behind an Application Load Balancer (ALB) in an Auto Scaling group. The site is experiencing performance issues related to a high request rate from illegitimate external systems with changing IP addresses. The security team is worried about potential DDoS attacks against the website. The company must block the illegitimate incoming requests in a way that has a minimal impact on legitimate users.

What should a solutions architect recommend?

- A. Deploy Amazon Inspector and associate it with the ALB.
- B. Deploy AWS WAF, associate it with the ALB, and configure a rate-limiting rule. **Most Voted**
- C. Deploy rules to the network ACLs associated with the ALB to block the incoming traffic.
- D. Deploy Amazon GuardDuty and enable rate-limiting protection when configuring GuardDuty.

Correct Answer: B

Community vote distribution



Comments

samehpalass **Highly Voted** 1 year, 5 months ago

Selected Answer: B

As no shield protect here so WAF rate limit
upvoted 11 times

hydro143 1 year, 2 months ago

Where's your Shield Advanced now, in your hour of need he has abandoned you
upvoted 11 times

pentium75 **Highly Voted** 11 months, 1 week ago

Selected Answer: B

Best solution Shield Advanced, not listed here, thus second-best solution, WAF with rate limiting
upvoted 8 times

awsgeek75 **Most Recent** 11 months ago

Selected Answer: B

A. Amazon Inspector = Software vulnerabilities like OS patches etc. Not fit for purpose.
C. Changing IP from DDoS so don't know the incoming traffic for configuration (even if it was possible)
D. GuardDuty is for workload and AWS account monitoring so it can't help with DDoS.

B is correct as AWS WAF + ALB can configure rate limiting even if source IP changes.

upvoted 6 times

jAtlas7 11 months, 2 weeks ago

Selected Answer: B

according to some google searches... to protect against DDOS attack:

* AWS WAF(Web Application Firewall) provides protection on the application layer (I think Application Load Balancer belongs to this level)

* AWS Shield protects the infrastructure layers of the OSI mode (I think AWS Network Load Balancer belongs to this level)

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

This case is A

upvoted 1 times

vincent2023 4 months, 2 weeks ago

Inspector wouldn't help in this case as it's used to detect vulnerabilities.

upvoted 1 times

pentium75 11 months, 1 week ago

Inspector is for detecting vulnerabilities, has nothing to do with the requirement.

upvoted 3 times

james2033 1 year, 4 months ago

Selected Answer: B

AWS Web Application Firewall (WAF) + ALB (Application Load Balancer) See image at <https://aws.amazon.com/waf/> .
<https://docs.aws.amazon.com/waf/latest/developerguide/ddos-responding.html> .

Question keyword "high request rate", answer keyword "rate-limiting rule"

<https://docs.aws.amazon.com/waf/latest/developerguide/waf-rate-based-example-limit-login-page-keys.html>

Amazon GuardDuty for threat detection <https://aws.amazon.com/guardduty/> , not for DDoS.

upvoted 3 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: B

B in swahili 'ba' :)

external systems, incoming requests = AWS WAF

upvoted 2 times

Axeashes 1 year, 5 months ago

Selected Answer: B

layer 7 DDoS protection with WAF

<https://docs.aws.amazon.com/waf/latest/developerguide/ddos-get-started-web-acl-rbr.html>

upvoted 2 times

antropaws 1 year, 6 months ago

Selected Answer: B

B no doubt.

upvoted 1 times

Joselucho38 1 year, 6 months ago

Selected Answer: B

AWS WAF (Web Application Firewall) is a service that provides protection for web applications against common web exploits. By associating AWS WAF with the Application Load Balancer (ALB), you can inspect incoming traffic and define rules to allow or block requests based on various criteria.

upvoted 5 times

cloudenthusiast 1 year, 6 months ago

B

AWS Web Application Firewall (WAF) is a service that helps protect web applications from common web exploits and provides advanced security features. By deploying AWS WAF and associating it with the ALB, the company can set up rules to filter and block incoming requests based on specific criteria, such as IP addresses.

In this scenario, the company is facing performance issues due to a high request rate from illegitimate external systems with changing IP addresses. By configuring a rate-limiting rule in AWS WAF, the company can restrict the number of requests coming from each IP address, preventing excessive traffic from overwhelming the website. This will help mitigate the impact of potential DDoS attacks and ensure that legitimate users can access the site without interruption.

upvoted 5 times

Efren 1 year, 6 months ago

Selected Answer: B

If not AWS Shield, then WAF

upvoted 4 times

nosense 1 year, 6 months ago

Selected Answer: B

B obv for this

upvoted 4 times

Efren 1 year, 6 months ago

My mind slipped with AWS Shield. GuardDuty can be working along with WAF for DDOS attack, but ultimately would be WAF

<https://aws.amazon.com/blogs/security/how-to-use-amazon-guardduty-and-aws-web-application-firewall-to-automatically-block-suspicious-hosts/>

upvoted 2 times

Mia2009687 1 year, 4 months ago

Same here, I was looking for AWS Shield

upvoted 1 times

Efren 1 year, 6 months ago

Selected Answer: D

D, Guard Duty for me

upvoted 1 times

pentium75 11 months, 1 week ago

Guard Duty detects threats, has nothing to do with rate-limiting.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #438

Topic 1

A company wants to share accounting data with an external auditor. The data is stored in an Amazon RDS DB instance that resides in a private subnet. The auditor has its own AWS account and requires its own copy of the database.

What is the MOST secure way for the company to share the database with the auditor?

- A. Create a read replica of the database. Configure IAM standard database authentication to grant the auditor access.
- B. Export the database contents to text files. Store the files in an Amazon S3 bucket. Create a new IAM user for the auditor. Grant the user access to the S3 bucket.
- C. Copy a snapshot of the database to an Amazon S3 bucket. Create an IAM user. Share the user's keys with the auditor to grant access to the object in the S3 bucket.
- D. Create an encrypted snapshot of the database. Share the snapshot with the auditor. Allow access to the AWS Key Management Service (AWS KMS) encryption key. **Most Voted**

Correct Answer: D*Community vote distribution*

D (100%)

Comments**alexandercamachop** **Highly Voted** 2 years ago**Selected Answer: D**

The most secure way for the company to share the database with the auditor is option D: Create an encrypted snapshot of the database, share the snapshot with the auditor, and allow access to the AWS Key Management Service (AWS KMS) encryption key.

By creating an encrypted snapshot, the company ensures that the database data is protected at rest. Sharing the encrypted snapshot with the auditor allows them to have their own copy of the database securely.

In addition, granting access to the AWS KMS encryption key ensures that the auditor has the necessary permissions to decrypt and access the encrypted snapshot. This allows the auditor to restore the snapshot and access the data securely.

This approach provides both data protection and access control, ensuring that the database is securely shared with the auditor while maintaining the confidentiality and integrity of the data.

upvoted 23 times

Total 1 vote 1 year 10 months ago

I anqkpkemei 1 year, 12 months ago

best explanation ever
upvoted 4 times

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: D

Key word: "Secure way"
The snapshot contents are encrypted using KMS keys for data security.
Sharing the snapshot directly removes risks of extracting/transferring data.
The auditor can restore the snapshot into their own RDS instance.
Access is controlled through sharing the encrypted snapshot and KMS key.
upvoted 5 times

24b2e9e **Most Recent** 11 months, 3 weeks ago

why not A ?
upvoted 5 times

LeonSauveterre 6 months, 1 week ago

I think it's simply because it introduces more risks than sharing just a snapshot. External entities don't need that many privileges, so it's not the MOST secure way
upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: D

Encrypted snapshot must be most secure compare others
upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_ShareSnapshot.html

With AWS RDS, you can share snapshots across accounts so no need to go through S3 or replication. Option D allows more secure way by using encryption and sharing encryption key.
upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: D

MOST secure way
upvoted 3 times

antropaws 2 years ago

Selected Answer: D

Most likely D.
upvoted 4 times

cloudenthusiast 2 years ago

Option D (Creating an encrypted snapshot of the database, sharing the snapshot, and allowing access to the AWS Key Management Service encryption key) is generally considered a better option for sharing the database with the auditor in terms of security and control.
upvoted 3 times

nosense 2 years ago

Selected Answer: D

D for me
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #439

Topic 1

A solutions architect configured a VPC that has a small range of IP addresses. The number of Amazon EC2 instances that are in the VPC is increasing, and there is an insufficient number of IP addresses for future workloads.

Which solution resolves this issue with the LEAST operational overhead?

- A. Add an additional IPv4 CIDR block to increase the number of IP addresses and create additional subnets in the VPC. Create new resources in the new subnets by using the new CIDR. **Most Voted**
- B. Create a second VPC with additional subnets. Use a peering connection to connect the second VPC with the first VPC. Update the routes and create new resources in the subnets of the second VPC.
- C. Use AWS Transit Gateway to add a transit gateway and connect a second VPC with the first VPC. Update the routes of the transit gateway and VPCs. Create new resources in the subnets of the second VPC.
- D. Create a second VPC. Create a Site-to-Site VPN connection between the first VPC and the second VPC by using a VPN-hosted solution on Amazon EC2 and a virtual private gateway. Update the route between VPCs to the traffic through the VPN. Create new resources in the subnets of the second VPC.

Correct Answer: A

Community vote distribution

A (100%)

Comments

antropaws **Highly Voted** 2 years ago

Selected Answer: A

A is correct: You assign a single CIDR IP address range as the primary CIDR block when you create a VPC and can add up to four secondary CIDR blocks after creation of the VPC.

upvoted 8 times

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: A

the architect just needs to:

Add the CIDR using the AWS console or CLI
Create new subnets in the VPC using the new CIDR

Launch resources in the new subnets

upvoted 5 times

Tendu Most Recent 3 months, 2 weeks ago

Selected Answer: A

A - is the correct answer

upvoted 1 times

rosanna 6 months ago

Selected Answer: A

The CIDR block of a VPC can NOT be changed but can be extended.

Thus, for the LEAST operational overhead; extend the existing one.

upvoted 1 times

f2e2419 1 year, 5 months ago

Selected Answer: A

best option

upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

A: LEAST operational overhead is by creating a new CIDR block in existing VPC.

All other options require additional overhead of gateway or second VPC

upvoted 4 times

potomac 1 year, 7 months ago

Selected Answer: A

After you've created your VPC, you can associate additional IPv4 CIDR blocks with the VPC

upvoted 4 times

TariqKipkemei 1 year, 12 months ago

Selected Answer: A

A is best

upvoted 2 times

Yadav_Sanjay 2 years ago

Selected Answer: A

Add additional CIDR of bigger range

upvoted 3 times

Efren 2 years ago

Selected Answer: A

Add new bigger subnets

upvoted 3 times

nosense 2 years ago

Selected Answer: A

A valid

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #440

Topic 1

A company used an Amazon RDS for MySQL DB instance during application testing. Before terminating the DB instance at the end of the test cycle, a solutions architect created two backups. The solutions architect created the first backup by using the mysqldump utility to create a database dump. The solutions architect created the second backup by enabling the final DB snapshot option on RDS termination.


The company is now planning for a new test cycle and wants to create a new DB instance from the most recent backup. The company has chosen a MySQL-compatible edition of Amazon Aurora to host the DB instance.

Which solutions will create the new DB instance? (Choose two.)

- A. Import the RDS snapshot directly into Aurora. **Most Voted**
- B. Upload the RDS snapshot to Amazon S3. Then import the RDS snapshot into Aurora.
- C. Upload the database dump to Amazon S3. Then import the database dump into Aurora. **Most Voted**
- D. Use AWS Database Migration Service (AWS DMS) to import the RDS snapshot into Aurora.
- E. Upload the database dump to Amazon S3. Then use AWS Database Migration Service (AWS DMS) to import the database dump into Aurora.

Correct Answer: AC

Community vote distribution

**Comments**

Axaus **Highly Voted** 2 years ago

Selected Answer: AC

A,C

A because the snapshot is already stored in AWS.

C because you don't need a migration tool going from MySQL to MySQL. You would use the MySQL utility.

upvoted 13 times

oras2023 **Highly Voted** 2 years ago

Selected Answer: AC

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/AuroraMySQL.Migrating.RDSMySQL.Import.html>
upvoted 8 times

bora4motion Most Recent 1 month, 1 week ago

Selected Answer: CE

c + e is correct
upvoted 1 times

tch 2 months, 2 weeks ago

Selected Answer: AC

Answer B is out, why upload to S3?
This is to call the most recent backup, answer D & E are out to use AWS Database Migration Service (AWS DMS).
upvoted 1 times

rpmaws 9 months ago

Selected Answer: CE

AWS DMS does not support migrating data directly from an RDS snapshot. DMS can migrate data from a live RDS instance or from a database dump, but not from a snapshot.
Also dump can be migrated to aurora using sql client.
upvoted 3 times

Mayur_B 7 months, 3 weeks ago

Correct.
Homogeneous migration in AWS Database Migration Service (DMS) refers to migrating databases between the same type of database engine. Here's an example:

Example:
Source Database: Amazon RDS for MySQL
Target Database: Amazon Aurora MySQL
upvoted 1 times

JackyCCK 1 year, 2 months ago

If you can use option A - "Import the RDS snapshot directly into Aurora", why go S3 in option C ?
Non sense, A and C cannot co-exist
upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: AC

A per <https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/AuroraMySQL.Migrating.RDSMySQL.Snapshot.html>

C per <https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/AuroraMySQL.Migrating.ExtMySQL.html>
upvoted 7 times

aws94 1 year, 5 months ago

Selected Answer: AB

A and B
upvoted 1 times

meowruki 1 year, 6 months ago

Selected Answer: AC

Similar : <https://repost.aws/knowledge-center/aurora-postgresql-migrate-from-rds>
upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: AD

A and C
upvoted 1 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: AC

Either import the RDS snapshot directly into Aurora or upload the database dump to Amazon S3, then import the database dump into Aurora.

upvoted 3 times

thanhnv142 1 year, 7 months ago

AC:

- store dump in s3 then upload to aurora
- no need to store snapshot in s3 because is in AWS already

upvoted 5 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: CE

C and E are the solutions that can restore the backups into Amazon Aurora.

The RDS DB snapshot contains backup data in a proprietary format that cannot be directly imported into Aurora. The mysqldump database dump contains SQL statements that can be imported into Aurora after uploading to S3. AWS DMS can migrate the dump file from S3 into Aurora.

upvoted 5 times

james2033 1 year, 10 months ago

Selected Answer: AC

Amazon RDS for MySQL --> Amazon Aurora MySQL-compatible.

* mysqldump, database dump --> (C) Upload to Amazon S3, Import dump to Aurora.

* DB snapshot --> (A) Import RDS Snapshot directly Aurora. The correct word should be "migration". "Use console to migrate the DB snapshot and create an Aurora MySQL DB cluster with the same databases as the original MySQL DB instance."

Exclude B, because no need upload DB snapshot to Amazon S3. Exclude D, because no need Migration service. Exclude E, because no need Migration service. Use exclusion method is more easy for this question.

Related links:

- Amazon RDS create database snapshot

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_CreateSnapshot.html

- <https://aws.amazon.com/rds/aurora/>

upvoted 3 times

marufxplorer 1 year, 11 months ago

CE

Since the backup created by the solutions architect was a database dump using the mysqldump utility, it cannot be directly imported into Aurora using RDS snapshots. Amazon Aurora has its own specific backup format that is different from RDS snapshots

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

C and E are the solutions that can restore the backups into Amazon Aurora.

The RDS DB snapshot contains backup data in a proprietary format that cannot be directly imported into Aurora. The mysqldump database dump contains SQL statements that can be imported into Aurora after uploading to S3. AWS DMS can migrate the dump file from S3 into Aurora.

upvoted 3 times

antropaws 2 years ago

Selected Answer: AC

Migrating data from MySQL by using an Amazon S3 bucket

You can copy the full and incremental backup files from your source MySQL version 5.7 database to an Amazon S3 bucket, and then restore an Amazon Aurora MySQL DB cluster from those files.

This option can be considerably faster than migrating data using mysqldump, because using mysqldump replays all of the commands to recreate the schema and data from your source database in your new Aurora MySQL DB cluster.

By copying your source MySQL data files, Aurora MySQL can immediately use those files as the data for an Aurora MySQL DB cluster

cluster.

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/AuroraMySQL.Migrating.ExtMySQL.html>

upvoted 4 times

omoakin 2 years ago

BE

Upload the RDS snapshot to Amazon S3. Then import the RDS snapshot into Aurora.

Upload the database dump to Amazon S3. Then use AWS Database Migration Service (AWS DMS) to import the database dump into Aurora

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #441

Topic 1

A company hosts a multi-tier web application on Amazon Linux Amazon EC2 instances behind an Application Load Balancer. The instances run in an Auto Scaling group across multiple Availability Zones. The company observes that the Auto Scaling group launches more On-Demand Instances when the application's end users access high volumes of static web content. The company wants to optimize cost.

What should a solutions architect do to redesign the application MOST cost-effectively?

- A. Update the Auto Scaling group to use Reserved Instances instead of On-Demand Instances.
- B. Update the Auto Scaling group to scale by launching Spot Instances instead of On-Demand Instances.
- C. Create an Amazon CloudFront distribution to host the static web contents from an Amazon S3 bucket. **Most Voted**
- D. Create an AWS Lambda function behind an Amazon API Gateway API to host the static website contents.

Correct Answer: C

Community vote distribution

C (100%)

Comments

awsgeek75 11 months ago

Selected Answer: C

C: Cost effective static content scaling = CloudFront
 A and B scale instances so not the best use of money for static content
 D Probably most expensive way of service static content at scale as you'll be charged for Lambda execution also
 upvoted 4 times

mwwt2022 11 months ago

Selected Answer: C

static content -> CloudFront
 upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

implementing CloudFront to serve static content is the most cost-optimal architectural change for this use case.

upvoted 4 times

james2033 1 year, 4 months ago

Selected Answer: C

Keyword "Amazon CloudFront", "high volumes of static web content", choose C.

upvoted 3 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: C

static web content = Amazon CloudFront

upvoted 2 times

alexandercamachop 1 year, 6 months ago

Selected Answer: C

Static Web Content = S3 Always.

CloudFront = Closer to the users locations since it will cache in the Edge nodes.

upvoted 3 times

cloudenthusiast 1 year, 6 months ago

By leveraging Amazon CloudFront, you can cache and serve the static web content from edge locations worldwide, reducing the load on your EC2 instances. This can help lower the number of On-Demand Instances required to handle high volumes of static web content requests. Storing the static content in an Amazon S3 bucket and using CloudFront as a content delivery network (CDN) improves performance and reduces costs by reducing the load on your EC2 instances.

upvoted 4 times

Efren 1 year, 6 months ago

Selected Answer: C

Static content, cloudFront plus S3

upvoted 3 times

nosense 1 year, 6 months ago

Selected Answer: C

c for me

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #442

Topic 1

A company stores several petabytes of data across multiple AWS accounts. The company uses AWS Lake Formation to manage its data lake. The company's data science team wants to securely share selective data from its accounts with the company's engineering team for analytical purposes.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Copy the required data to a common account. Create an IAM access role in that account. Grant access by specifying a permission policy that includes users from the engineering team accounts as trusted entities.
- B. Use the Lake Formation permissions Grant command in each account where the data is stored to allow the required engineering team users to access the data.
- C. Use AWS Data Exchange to privately publish the required data to the required engineering team accounts.
- D. Use Lake Formation tag-based access control to authorize and grant cross-account permissions for the required data to the engineering team accounts. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

cloudenthusiast **Highly Voted** 1 year, 6 months ago

Selected Answer: D

By utilizing Lake Formation's tag-based access control, you can define tags and tag-based policies to grant selective access to the required data for the engineering team accounts. This approach allows you to control access at a granular level without the need to copy or move the data to a common account or manage permissions individually in each account. It provides a centralized and scalable solution for securely sharing data across accounts with minimal operational overhead.

upvoted 18 times

NSA_Poker **Most Recent** 6 months, 1 week ago

Selected Answer: D

(B) uses the CLI command that has many options: principal, TableName, ColumnNames, LFTag etc providing a way to manage granular access permissions for different users at the table and column level. That way you don't give full access to the all the data. The problem with (B) is to implement this in each account has a lot more operational overhead than (D).

upvoted 2 times

upvoted 2 times

awsgeek75 11 months ago

Selected Answer: D

D: Selective data = tagging

A and B gives full access to all the data

C is possible but with complex operational overhead as you have to publish your data to the Data Exchange. (this is based on my limited knowledge so happy to be corrected)

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

D is the correct option with the least operational overhead.

Using Lake Formation tag-based access control allows granting cross-account permissions to access data in other accounts based on tags, without having to copy data or configure individual permissions in each account.

This provides a centralized, tag-based way to share selective data across accounts to authorized users with least operational overhead.

upvoted 4 times

luisgu 1 year, 6 months ago

Selected Answer: D

<https://aws.amazon.com/blogs/big-data/securely-share-your-data-across-aws-accounts-using-aws-lake-formation/>

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #443

Topic 1


A company wants to host a scalable web application on AWS. The application will be accessed by users from different geographic regions of the world. Application users will be able to download and upload unique data up to gigabytes in size. The development team wants a cost-effective solution to minimize upload and download latency and maximize performance.

What should a solutions architect do to accomplish this?

- A. Use Amazon S3 with Transfer Acceleration to host the application. **Most Voted**
- B. Use Amazon S3 with CacheControl headers to host the application.
- C. Use Amazon EC2 with Auto Scaling and Amazon CloudFront to host the application.
- D. Use Amazon EC2 with Auto Scaling and Amazon ElastiCache to host the application.

Correct Answer: A

Community vote distribution



Comments

pentium75 **Highly Voted** 1 year, 5 months ago

Selected Answer: A

The question asks for "a cost-effective solution [ONLY TO] to minimize upload and download latency and maximize performance", not for the actual application. And the 'cost-effective solution to minimize upload and download latency and maximize performance' is S3 Transfer Acceleration. Obviously there is more required to host the app, but that is not asked for.
upvoted 15 times

chris0975 **Highly Voted** 1 year, 7 months ago

Selected Answer: A

The question is focused on large downloads and uploads. S3 Transfer Acceleration is what fits. CloudFront is for caching which cannot be used when the data is unique. They aren't as concerned with regular web traffic.

Amazon S3 Transfer Acceleration can speed up content transfers to and from Amazon S3 by as much as 50-500% for long-distance transfer of larger objects.

upvoted 7 times

bora4motion **Most Recent** 1 month, 1 week ago

Selected Answer: A

voting for A.
upvoted 1 times

Ben_88 7 months, 2 weeks ago

Selected Answer: C

S3 can't "host an application", so even tho S3 Transfer acceleration would optimize the performances of the download/upload, still it doesn't host the app. the question is poorly written i guess
upvoted 1 times

vincent2023 4 months, 2 weeks ago

I agree with your point !!
upvoted 1 times

ChymKuBoy 9 months, 3 weeks ago

Selected Answer: A

A for sure
upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

Not C, D No requirements to scale the application itself so EC2 is not applicable.
B is for caching so not sure how/if that helps the upload speed for global users
A is correct as Transfer Accelerator is best for uploading and downloading unique items near the user's region/location
upvoted 6 times

tosuccess 1 year, 5 months ago

Selected Answer: A

for datas greater than 1 GB, s3 transfer acceleration is the best
upvoted 3 times

Cyberkayu 1 year, 5 months ago

Selected Answer: A

Application users will be able to download and upload UNIQUE data up to gigabytes in size

Thus all caching related solution dont work.
upvoted 5 times

Goutham4981 1 year, 6 months ago

Selected Answer: A

Downloading data upto gigabytes in size - Cloudfront is a content delivery service that acts as an edge caching layer for images and other data. Not a service that minimizes upload and download latency.
upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: A

The question is focused on large downloads and uploads. S3 Transfer Acceleration is what fits. CloudFront is for caching which cannot be used when the data is unique. They aren't as concerned with regular web traffic.

C didn't mention S3. Where the data is stored?
upvoted 4 times

pentium75 1 year, 5 months ago

A doesn't mention EC2 or EKS or ECS or Elastic Beanstalk or Lambda. Where does the "scalable web application" run?
upvoted 2 times

JA2018 6 months, 2 weeks ago

hmm... could be a red herring?

upvoted 1 times

beast2091 1 year, 7 months ago

It is A.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/transfer-acceleration.html>

upvoted 2 times

danielmakita 1 year, 7 months ago

It is A as the Transfer Acceleration will minimize upload and download latency.

If you choose C, where would the files be stored? There is no mention of any S3. Will it be stored inside the EC2? That's why I didn't go for C

upvoted 5 times

Sindokuhlep 1 year, 7 months ago

Selected Answer: C

Amazon S3 with Transfer Acceleration (option A) is designed for speeding up uploads to Amazon S3, and it's not used for hosting scalable web applications. It doesn't mention using EC2 instances for hosting the application.

upvoted 4 times

canonlycontainletters1 1 year, 7 months ago

Selected Answer: C

My answer is C

upvoted 1 times

thanhnv142 1 year, 7 months ago

C because A is for upload data to S3, not for web app

upvoted 1 times

DamyanG 1 year, 8 months ago

Selected Answer: C

The correct answer is C!!! It is not A, because

- Amazon S3 with Transfer Acceleration (option A) is designed for speeding up uploads to Amazon S3, and it's not used for hosting scalable web applications. It doesn't mention using EC2 instances for hosting the application.

upvoted 3 times

Victory007 1 year, 8 months ago

Selected Answer: C

Amazon CloudFront is a global content delivery network (CDN) that delivers web content to users with low latency and high transfer speeds. It does this by caching content at edge locations around the world, which are closer to the users than the origin server.

By using Amazon EC2 with Auto Scaling and Amazon CloudFront, the company can create a scalable and high-performance web application that is accessible to users from different geographic regions of the world.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #444

Topic 1

A company has hired a solutions architect to design a reliable architecture for its application. The application consists of one Amazon RDS DB instance and two manually provisioned Amazon EC2 instances that run web servers. The EC2 instances are located in a single Availability Zone.

An employee recently deleted the DB instance, and the application was unavailable for 24 hours as a result. The company is concerned with the overall reliability of its environment.

What should the solutions architect do to maximize reliability of the application's infrastructure?

- A. Delete one EC2 instance and enable termination protection on the other EC2 instance. Update the DB instance to be Multi-AZ, and enable deletion protection.
- B. Update the DB instance to be Multi-AZ, and enable deletion protection. Place the EC2 instances behind an Application Load Balancer, and run them in an EC2 Auto Scaling group across multiple Availability Zones. **Most Voted**
- C. Create an additional DB instance along with an Amazon API Gateway and an AWS Lambda function. Configure the application to invoke the Lambda function through API Gateway. Have the Lambda function write the data to the two DB instances.
- D. Place the EC2 instances in an EC2 Auto Scaling group that has multiple subnets located in multiple Availability Zones. Use Spot Instances instead of On-Demand Instances. Set up Amazon CloudWatch alarms to monitor the health of the instances. Update the DB instance to be Multi-AZ, and enable deletion protection.

Correct Answer: B

Community vote distribution

B (100%)

Comments

awsgeek75 **Highly Voted** 11 months ago

Option E: Sack the employee who did this :)
upvoted 10 times

nosense **Highly Voted** 1 year, 6 months ago

Selected Answer: B

B is correct. HA ensured by DB in Mutli-AZ and EC2 in AG

upvoted 5 times

wizcloudifa Most Recent 7 months, 3 weeks ago

Selected Answer: B

A: delete one instance, why?. Although takes care of reliability of DB instance however not EC2.

B. seems perfect as takes care of reliability of both EC2 as well as DB

C. DB instance's reliability is not taken care of

D. seems to be trying to address cost alongside reliability of EC2 and DB.

upvoted 3 times

awsgeek75 11 months ago

Selected Answer: B

A: Deleting one EC2 instance makes no sense. Why would you do that?

C: API Gateway, Lambda etc are all nice but they don't solve the problem of DB instance deletion

D: EC2 subnet blah blah, what? The problem is reliability, not networking!

B is correct as it solves the DB deletion issue and increases reliability by Multi AZ scaling of EC2 instances

upvoted 5 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

The key points:

° RDS Multi-AZ and deletion protection provide high availability for the database.

° The load balancer and Auto Scaling group across AZs give high availability for EC2.

° Options A, C, D have limitations that would reduce reliability vs option B.

upvoted 3 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: B

Update the DB instance to be Multi-AZ, and enable deletion protection. Place the EC2 instances behind an Application Load Balancer, and run them in an EC2 Auto Scaling group across multiple Availability Zones

upvoted 2 times

antropaws 1 year, 6 months ago

Selected Answer: B

B for sure.

upvoted 2 times

alexandercamachop 1 year, 6 months ago

Selected Answer: B

It is the only one with High Availability.

Amazon RDS with Multi AZ

EC2 with Auto Scaling Group in Multi Az

upvoted 2 times

omoakin 1 year, 6 months ago

same question from

<https://www.examttopics.com/exams/amazon/aws-certified-solutions-architect-associate-saa-c02/>

long time ago and still same option B

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #445

Topic 1

A company is storing 700 terabytes of data on a large network-attached storage (NAS) system in its corporate data center. The company has a hybrid environment with a 10 Gbps AWS Direct Connect connection.

After an audit from a regulator, the company has 90 days to move the data to the cloud. The company needs to move the data efficiently and without disruption. The company still needs to be able to access and update the data during the transfer window.

Which solution will meet these requirements?

- A. Create an AWS DataSync agent in the corporate data center. Create a data transfer task Start the transfer to an Amazon S3 bucket. **Most Voted**
- B. Back up the data to AWS Snowball Edge Storage Optimized devices. Ship the devices to an AWS data center. Mount a target Amazon S3 bucket on the on-premises file system.
- C. Use rsync to copy the data directly from local storage to a designated Amazon S3 bucket over the Direct Connect connection.
- D. Back up the data on tapes. Ship the tapes to an AWS data center. Mount a target Amazon S3 bucket on the on-premises file system.

Correct Answer: A

Community vote distribution

A (100%)

Comments

wRhIH **Highly Voted** 1 year, 12 months ago

For those who wonders why not B. Snowball Edge Storage Optimized device for data transfer is up to 100TB
<https://docs.aws.amazon.com/snowball/latest/developer-guide/device-differences.html>
upvoted 12 times

Maru86 1 year, 2 months ago

The question explicitly mentioned "devices", also Snowball Edge Storage Optimized is 80TB HDD. So it is possible, but the answer is A because we can transfer with DataSync in 6.5 days.

upvoted 4 times

smartegnine 1 year, 11 months ago

$10\text{GBs} * 24 * 60 * 60 = 864,000 \text{ GB}$ estimate around 864 TB a day, 2 days will transfer all data. But for snowball at least 4 days for delivery to the data center.

upvoted 1 times

siGma182 1 year, 10 months ago

This account is wrong but I get your point. It is wrong cause 10Gb/s is not the same as 10GB/s (Gigabits vs Gigabytes). However, the correct count is $864\text{Tb}/8 = 108\text{TB}$ per day. In one week you should've transferred all the data.

upvoted 8 times

JA2018 6 months, 2 weeks ago

folks, you had to consider the typical TCP overheads for such over-the-wire data transfers.... it could be as high as 70% of the total available bandwidth due to a wide variety of factors

upvoted 1 times

Maru86 1 year, 2 months ago

That's right, 1 GB = 8 Gb. Essentially we have a speed of 1.25GB/s.

upvoted 2 times

hsinchang Highly Voted 1 year, 10 months ago

Selected Answer: A

Access during the transfer window -> DataSync

upvoted 9 times

surajkrishnamurthy Most Recent 4 months ago

Selected Answer: A

700TB with a Bandwidth of 10Gbps takes approximately 7 days, 3 Hours, 2 Minutes, 7 Seconds

As per the below link>>

<https://expedient.com/knowledgebase/tools-and-calculators/file-transfer-time-calculator/>

So Option A is Correct & using AWS Datasync facilitates further access to data still present in the OnPrem to access and S3 provides with unlimited data storage capacity

upvoted 1 times

Danilus 7 months ago

A- Porque La compañía todavía necesita poder acceder y actualizar los datos durante la ventana de transferencia

B-Aunque esta podría ser una opción porque dice de enviar los dispositivos al centro de datos que serían más o menos entre 9 y 10 dispositivos ya que el almacenamiento de storage optimized es de 80 TB el problema está en que con Snowball Edge Storage Optimized no se puede acceder y actualizar a los datos durante la transferencia

upvoted 1 times

JA2018 6 months, 2 weeks ago

From Google Translate:

A- Because The company still needs to be able to access and update the data during the transfer window

B-Although this could be an option because it says that sending the devices to the data center would be more or less between 9 and 10 devices since the storage optimized storage is 80 TB, the problem is that with Snowball Edge Storage Optimized it does not Data can be accessed and updated during transfer

upvoted 3 times

Danilus 7 months ago

Selected Answer: A

A- Porque La compañía todavía necesita poder acceder y actualizar los datos durante la ventana de transferencia

B-Aunque esta podría ser una opción porque dice de enviar los dispositivos al centro de datos que serían más o menos entre 9 y 10 dispositivos ya que el almacenamiento de storage optimized es de 80 TB el problema está en que con Snowball Edge Storage Optimized no se puede acceder y actualizar a los datos durante la transferencia

upvoted 1 times

MetAlus 8 months, 4 weeks ago

maldives 6 months, 4 weeks ago

Selected Answer: A

Finally, a company with good bandwidth.
upvoted 4 times

NSA_Poker 1 year ago

Selected Answer: A

(B) is incorrect bc although Mountpoint for S3 is possible for on-premises NAS, this is not as efficient as AWS DataSync. Data updates made during the transfer window would have to be resolved later.

upvoted 2 times

Burrito69 1 year, 2 months ago

I will put simple calculation for oyu guys to just store in your head to quickly answer:
for 10GBPS its 1.25GBPS becasue its bits to bytes.
for one minute its 75GBPS
for one hour its 4500 GBPS
for one day its 10.8 TB

so you can calculate easily if you just store these numbers in your head. lets say if question is 1GBPS DirectConnect that meand everything above should be divide by 8. cool

upvoted 1 times

JA2018 6 months, 2 weeks ago

folks, you had to consider the typical TCP overheads for such over-the-wire data transfers.... it could be as high as 70% of the total available bandwidth due to a wide variety of factors

in a worst case scenario, we could be looking @ 3.24 TB/ day

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

Critical requirement: "The company needs to move the data efficiently and without disruption."

B: Causes disruption

C: I don't think that is possible without a gateway kind of thing

D: Tape backups? " Mount a target Amazon S3 bucket on the on-premises file system"? This requires some gateway which is not mentioned

A is the answer as DataSync allows transfer without disruption and with 10Gbps, it can be done in 90 days.

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

AWS DataSync can efficiently transfer large datasets from on-premises NAS to Amazon S3 over Direct Connect.

DataSync allows accessing and updating the data continuously during the transfer process.

upvoted 5 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: A

AWS DataSync is a secure, online service that automates and accelerates moving data between on premises and AWS Storage services.

upvoted 3 times

omoakin 2 years ago

A

<https://www.examttopics.com/discussions/amazon/view/46492-exam-aws-certified-solutions-architect-associate-saa-c02/#:~:text=Exam%20question%20from,Question%20%23%3A%20385>

upvoted 2 times

cloudenthusiast 2 years ago

Selected Answer: A

By leveraging AWS DataSync in combination with AWS Direct Connect, the company can efficiently and securely transfer its 700 terabytes of data to an Amazon S3 bucket without disruption. The solution allows continued access and updates to the data during the transfer window, ensuring business continuity throughout the migration process.

upvoted 4 times

nosense 2 years ago

Selected Answer: A

A for me, bcs egde storage up to 100tb

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #446

Topic 1


A company stores data in PDF format in an Amazon S3 bucket. The company must follow a legal requirement to retain all new and existing data in Amazon S3 for 7 years.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Turn on the S3 Versioning feature for the S3 bucket. Configure S3 Lifecycle to delete the data after 7 years. Configure multi-factor authentication (MFA) delete for all S3 objects.
- B. Turn on S3 Object Lock with governance retention mode for the S3 bucket. Set the retention period to expire after 7 years. Recopy all existing objects to bring the existing data into compliance.
- C. Turn on S3 Object Lock with compliance retention mode for the S3 bucket. Set the retention period to expire after 7 years. Recopy all existing objects to bring the existing data into compliance.
- D. Turn on S3 Object Lock with compliance retention mode for the S3 bucket. Set the retention period to expire after 7 years. Use S3 Batch Operations to bring the existing data into compliance. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

omarshaban **Highly Voted** 1 year, 4 months ago

THIS WAS IN MY EXAM
upvoted 11 times

awsgeek75 1 year, 4 months ago

Did you pass?
upvoted 4 times

DrWatson **Highly Voted** 2 years ago

Selected Answer: D

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/batch-ops-retention-date.html>
upvoted 6 times

Lin878 Most Recent 11 months, 3 weeks ago

Selected Answer: D

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock.html>
upvoted 3 times

iapps369 1 year, 4 months ago

D
as S3 batch operations reduce risk and manual copy/paste overhead.
upvoted 6 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

A: Versioning, not relevant
B: Governance, it won't enforce object lock
C: Recopy existing objects may work but lots of operational overhead (see link)
D: Compliance on existing objects with batch operations is least operational overhead

<https://repost.aws/questions/QUGKrl8XRLTEeulzUHq0lkew/s3-object-lock-on-existing-s3-objects>
upvoted 5 times

awsgeek75 1 year, 5 months ago

With option C, you have to copy the object for it to be compliant and then delete the original as only the new copy will be compliant. So D is the only option
upvoted 2 times

mr123dd 1 year, 5 months ago

Selected Answer: A

To enable Object Lock on an Amazon S3 bucket, you must first enable versioning on that bucket. other 3 option did not enable versioning first
upvoted 2 times

fb4afde 1 year, 5 months ago

Selected Answer: D

Recopying offers more control but requires users to manage the process. S3 Batch Operations automates the process at scale but with less granular control - LEAST operational overhead
upvoted 3 times

moonster 1 year, 6 months ago

Its C because you only need to recopy all existing objects one time, so why use S3 batch operations if new data's going to be in compliance retention mode? I can see why its C although my initial gut answer was D.
upvoted 2 times

pentium75 1 year, 5 months ago

What if I don't have the original files anymore? Where should I copy them from?
upvoted 3 times

kwang312 1 year, 8 months ago

You can only enable Object Lock for new buckets. If you want to turn on Object Lock for an existing bucket, contact AWS Support.
upvoted 2 times

pentium75 1 year, 5 months ago

You need a token from AWS Support, but you CAN enable Object Lock for an existing bucket.
upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

Turn on S3 Object Lock with compliance retention mode for the S3 bucket. Set the retention period to expire after 7 years. Use S3 Batch Operations to bring the existing data into compliance.

upvoted 2 times

MrAWSAssociate 1 year, 11 months ago

Selected Answer: D

To replicate existing object/data in S3 Bucket to bring them to compliance, optionally we use "S3 Batch Replication", so option D is the most appropriate, especially if we have big data in S3.

upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: D

For minimum ops D is best

upvoted 2 times

antropaws 2 years ago

Selected Answer: C

Batch operations will add operational overhead.

upvoted 3 times

pentium75 1 year, 5 months ago

And gathering all the files for copying them again does not?

upvoted 2 times

Abrar2022 2 years ago

Use Object Lock in Compliance mode. Then Use Batch operation.

WRONG>>manual work and not automated>>>Recopy all existing objects to bring the existing data into compliance.

upvoted 1 times

pentium75 1 year, 5 months ago

Batch IS automated. You just need to create the batch which is a one-time operation.

"Recopy all existing objects" is not operational overhead?

upvoted 2 times

omoakin 2 years ago

C

When an object is locked in compliance mode, its retention mode can't be changed, and its retention period can't be shortened. Compliance mode helps ensure that an object version can't be overwritten or deleted for the duration of the retention period.

upvoted 3 times

omoakin 2 years ago

error i meant to type D

i wont do recopy

upvoted 2 times

lucdt4 2 years ago

No, D for me because the requirement is LEAST operational overhead

So RECOpy is the manual operation -> C is wrong

D is correct

upvoted 3 times

cloudenthusiast 2 years ago

Recopying vs. S3 Batch Operations: In Option C, the recommendation is to recopy all existing objects to ensure they have the appropriate retention settings. This can be done using simple S3 copy operations. On the other hand, Option D suggests using S3 Batch Operations, which is a more advanced feature and may require additional configuration and management. S3 Batch Operations can be beneficial if you have a massive number of objects and need to perform complex operations, but it might introduce more overhead for this specific use case.

Operational complexity: Option C has a straightforward process of recopying existing objects. It is a well-known operation in S3 and doesn't require additional setup or management. Option D introduces the need to set up and configure S3 Batch Operations, which can involve creating job definitions, specifying job parameters, and monitoring the progress of batch operations. This additional complexity may increase the operational overhead.

upvoted 2 times

Efren 2 years ago

Selected Answer: D

You need AWS Batch to re-apply certain config to files that were already in S3, like encryption

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #447

Topic 1


A company has a stateless web application that runs on AWS Lambda functions that are invoked by Amazon API Gateway. The company wants to deploy the application across multiple AWS Regions to provide Regional failover capabilities.

What should a solutions architect do to route traffic to multiple Regions?

- A. Create Amazon Route 53 health checks for each Region. Use an active-active failover configuration. **Most Voted**
- B. Create an Amazon CloudFront distribution with an origin for each Region. Use CloudFront health checks to route traffic.
- C. Create a transit gateway. Attach the transit gateway to the API Gateway endpoint in each Region. Configure the transit gateway to route requests.
- D. Create an Application Load Balancer in the primary Region. Set the target group to point to the API Gateway endpoint hostnames in each Region.

Correct Answer: A

Community vote distribution



Comments

TariqKipkemei **Highly Voted** 1 year, 11 months ago

Selected Answer: A

Global, Reduce latency, health checks, no failover = Amazon CloudFront
 Global, Reduce latency, health checks, failover, Route traffic = Amazon Route 53
 option A has more weight.
 upvoted 33 times

ManikRoy 1 year ago

Cloud front does have failover capabilities.
https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/high_availability_origin_failover.html#:~:text=the%20secondary%20origin.-,Note,Choose%20Create%20origin%20group.
 upvoted 4 times

Anmol_1010 1 year, 7 months ago

nicley explained

upvoted 3 times

examtopictempacc Highly Voted 2 years ago

Selected Answer: A

A. I'm not an expert in this area, but I still want to express my opinion. After carefully reviewing the question and thinking about it for a long time, I actually don't know the reason. As I mentioned at the beginning, I'm not an expert in this field.

upvoted 19 times

awsgeek75 1 year, 4 months ago

All the explanation you need for this question and option A is in this article:

<https://aws.amazon.com/blogs/compute/building-a-multi-region-serverless-application-with-amazon-api-gateway-and-aws-lambda/>

upvoted 3 times

MatAlves Most Recent 8 months, 4 weeks ago

Selected Answer: A

Correct me if I'm wrong but CloudFront DOES NOT have health check capabilities out of the box. Route 53 and Global Accelerator do.

upvoted 3 times

ChymKuBoy 9 months, 2 weeks ago

Selected Answer: A

A for sure

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

B: Caching solution. Not ideal for failover although it will work. Would have been a correct answer if A wasn't an option

C: Transit gateway is for VPC connectivity not AWS API or Lambda

D: Even if it was possible, there is a primary region dependency of ALB

A: correct because R53 health checks can failover across regions

Good explanation here:

<https://aws.amazon.com/blogs/compute/building-a-multi-region-serverless-application-with-amazon-api-gateway-and-aws-lambda/>

upvoted 4 times

awsgeek75 1 year, 4 months ago

The article also explains why you cannot use a CloudFront distribution for API Gateway, Lambda for failover

upvoted 2 times

tosuccess 1 year, 5 months ago

Selected Answer: B

we can set primary and secondary regions in cloud front for failover.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: A

Application is serverless, it doesn't matter where it runs, so can be active-active setup and run wherever the request comes in. Route 53 with health checks will route to a healthy region.

B, could work too, but CloudFront is for caching which does not seem to help with an API. The goal here is "failover capabilities", not caching/performance/latency etc.

upvoted 3 times

Goutham4981 1 year, 6 months ago

Selected Answer: A

In active active failover config, route53 continuously monitors its endpoints and if one of them is unhealthy, it excludes the region/endpoint from its valid traffic route - Only Sensible option

Cloudfront is a content delivery network - not used to route traffic
Transit gateway for traffic routing - aws devs will hit us with a stick on hearing this option
You cant use a load balancer for cross region load balancing - invalid
upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: A

Global ,Reduce latency, health checks, failover, Route traffic = Amazon Route 53
upvoted 2 times

youdelin 1 year, 7 months ago

"What the?" yeah I know right
upvoted 1 times

jrestrepob 1 year, 9 months ago

Selected Answer: B

"Stateless applications provide one service or function and use content delivery network (CDN), web, or print servers to process these short-term requests.
<https://docs.aws.amazon.com/architecture-diagrams/latest/multi-region-api-gateway-with-cloudfront/multi-region-api-gateway-with-cloudfront.html>
upvoted 1 times

deechean 1 year, 9 months ago

its not static content, actually they deployed a API Gateway backed by lambda
upvoted 3 times

MrAWSAssociate 1 year, 11 months ago

Selected Answer: A

A option does make sense.
upvoted 2 times

Sangsation 1 year, 11 months ago

Selected Answer: B

By creating an Amazon CloudFront distribution with origins in each AWS Region where the application is deployed, you can leverage CloudFront's global edge network to route traffic to the closest available Region. CloudFront will automatically route the traffic based on the client's location and the health of the origins using CloudFront health checks.

Option A (creating Amazon Route 53 health checks with an active-active failover configuration) is not suitable for this scenario as it is primarily used for failover between different endpoints within the same Region, rather than routing traffic to different Regions.
upvoted 2 times

pentium75 1 year, 5 months ago

Option A does not speak of Route 53 failover routing policies.
upvoted 1 times

Axeashes 1 year, 12 months ago

Selected Answer: A

<https://aws.amazon.com/blogs/compute/building-a-multi-region-serverless-application-with-amazon-api-gateway-and-aws-lambda/>
upvoted 4 times

Gooniegoogoo 1 year, 11 months ago

that is from 2017.. i wonder if it is still relevant..
upvoted 1 times

DrWatson 2 years ago

Selected Answer: A

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/dns-failover.html>

upvoted 2 times

antropaws 2 years ago

Selected Answer: A

I understand that you can use Route 53 to provide regional failover.

upvoted 1 times

alexandercamachop 2 years ago

Selected Answer: A

To route traffic to multiple AWS Regions and provide regional failover capabilities for a stateless web application running on AWS Lambda functions invoked by Amazon API Gateway, you can use Amazon Route 53 with an active-active failover configuration.

By creating Amazon Route 53 health checks for each Region and configuring an active-active failover configuration, Route 53 can monitor the health of the endpoints in each Region and route traffic to healthy endpoints. In the event of a failure in one Region, Route 53 automatically routes traffic to the healthy endpoints in other Regions.

This setup ensures high availability and failover capabilities for your web application across multiple AWS Regions.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #448

Topic 1


A company has two VPCs named Management and Production. The Management VPC uses VPNs through a customer gateway to connect to a single device in the data center. The Production VPC uses a virtual private gateway with two attached AWS Direct Connect connections. The Management and Production VPCs both use a single VPC peering connection to allow communication between the applications.

What should a solutions architect do to mitigate any single point of failure in this architecture?

- A. Add a set of VPNs between the Management and Production VPCs.
- B. Add a second virtual private gateway and attach it to the Management VPC.
- C. Add a second set of VPNs to the Management VPC from a second customer gateway device. **Most Voted**
- D. Add a second VPC peering connection between the Management VPC and the Production VPC.

Correct Answer: C

Community vote distribution



Comments

Guru4Cloud **Highly Voted** 9 months, 2 weeks ago

Selected Answer: C

C is the correct option to mitigate the single point of failure.

The Management VPC currently has a single VPN connection through one customer gateway device. This is a single point of failure.

Adding a second set of VPN connections from the Management VPC to a second customer gateway device provides redundancy and eliminates this single point of failure.

upvoted 7 times

Guru4Cloud 9 months, 2 weeks ago

As @Abrar2022 explains
 (production) VPN 1-----> cgw 1
 (management) VPN 2-----> cgw
 upvoted 4 times

Abrar2022 Highly Voted 1 year ago

(production) VPN 1-----> cgw 1
(management) VPN 2-----> cgw 2
upvoted 6 times

Abrar2022 1 year ago

ANSWER IS C
upvoted 2 times

Dantecito Most Recent 4 months ago

Selected Answer: D

but what if the VPC peering connection fails? how The Management and Production vpc communicate?
upvoted 1 times

bsbs1234 8 months, 1 week ago

C,

(production) --PrivateGateway----->Direct Connect Gateway 1 ---> cgw 1 ---> DataCenter
(production) -- PrivateGateway -----> Direct Connect Gateway 2 --->cgw 2 --> DataCenter
(Management) -- > VPN ----> (Direct Connect Gateway 1?) --->cgw1 ----> dataCenter---> device in dataCenter
upvoted 2 times

omoakin 1 year ago

I agree to C
upvoted 2 times

cloudenthusiast 1 year ago

Selected Answer: C

option D is not a valid solution for mitigating single points of failure in the architecture. I apologize for the confusion caused by the incorrect information.

To mitigate single points of failure in the architecture, you can consider implementing option C: adding a second set of VPNs to the Management VPC from a second customer gateway device. This will introduce redundancy at the VPN connection level for the Management VPC, ensuring that if one customer gateway or VPN connection fails, the other connection can still provide connectivity to the data center.

upvoted 5 times

Efren 1 year ago

Selected Answer: C

Redundant VPN connections: Instead of relying on a single device in the data center, the Management VPC should have redundant VPN connections established through multiple customer gateways. This will ensure high availability and fault tolerance in case one of the VPN connections or customer gateways fails.

upvoted 5 times

nosense 1 year ago

Selected Answer: C

<https://www.examttopics.com/discussions/amazon/view/53908-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #449

Topic 1


A company runs its application on an Oracle database. The company plans to quickly migrate to AWS because of limited resources for the database, backup administration, and data center maintenance. The application uses third-party database features that require privileged access.

Which solution will help the company migrate the database to AWS MOST cost-effectively?

- A. Migrate the database to Amazon RDS for Oracle. Replace third-party features with cloud services.
- B. Migrate the database to Amazon RDS Custom for Oracle. Customize the database settings to support third-party features. **Most Voted**
- C. Migrate the database to an Amazon EC2 Amazon Machine Image (AMI) for Oracle. Customize the database settings to support third-party features.
- D. Migrate the database to Amazon RDS for PostgreSQL by rewriting the application code to remove dependency on Oracle APEX.

Correct Answer: B

Community vote distribution



Comments

Abrar2022 **Highly Voted** 1 year, 6 months ago

Selected Answer: B

RDS Custom since it's related to 3rd vendor
 RDS Custom since it's related to 3rd vendor
 RDS Custom since it's related to 3rd vendor
 upvoted 6 times

awsgeek75 **Most Recent** 11 months ago

Selected Answer: B

Key constraints: Limited resources for DB admin and cost. 3rd party db features with privileged access.

A: Won't work due to 3rd party features

C: AMI with Oracle may work but again overhead of backed, maintenance etc

D: Too much overhead in rewrite

B: Actually supports Oracle 3rd party features

B. Actually supports Oracle 3rd party features

Caution: If this is only about APEX as suggested in option D, then A is also a possible answer:

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.Oracle.Options.APEX.html>

upvoted 3 times

awsgeek75 11 months ago

Action ignore the last line of my previous comment, A is not a valid option in any case as it suggest replacing 3rd party features with cloud services which is not possible without more details.

upvoted 2 times

pentium75 11 months, 1 week ago

Selected Answer: B

"Amazon RDS Custom is a managed database service for applications that require customization of the underlying operating system and database environment. Benefits of RDS automation with the access needed for legacy, packaged, and custom applications."

That should allow the "privileged access".

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Migrate the database to Amazon RDS Custom for Oracle. Customize the database settings to support third-party features.

upvoted 4 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: B

Custom database features = Amazon RDS Custom for Oracle

upvoted 4 times

antropaws 1 year, 6 months ago

Selected Answer: B

Most likely B.

upvoted 2 times

omoakin 1 year, 6 months ago

CCCCCCCCCCCCCCCCCCCC

upvoted 1 times

aqmdla2002 1 year, 6 months ago

Selected Answer: B

<https://aws.amazon.com/about-aws/whats-new/2021/10/amazon-rds-custom-oracle/>

upvoted 3 times

hiroohiroo 1 year, 6 months ago

Selected Answer: B

https://docs.aws.amazon.com/ja_jp/AmazonRDS/latest/UserGuide/Oracle.Resources.html

upvoted 2 times

karbob 1 year, 6 months ago

Amazon RDS Custom for Oracle, which is not an actual service. !!!!

upvoted 2 times

nosense 1 year, 6 months ago

Option C is also a valid solution, but it is not as cost-effective as option B.

Option C requires the company to manage its own database infrastructure, which can be expensive and time-consuming. Additionally, the company will need to purchase and maintain Oracle licenses.

upvoted 3 times

v0 1 year, 6 months ago

RDS Custom enables the capability to access the underlying database and OS so as to configure additional settings to support 3rd party. This feature is applicable only for Oracle and Postgresql

upvoted 2 times

y0 1 year, 6 months ago

Sorry Oracle and sql server (not postgresql)

upvoted 2 times

omoakin 1 year, 6 months ago

I will say C cos of this

"application uses third-party "

upvoted 1 times

cloudenthusiast 1 year, 6 months ago

Selected Answer: C

Should not it be since for Ec2, the company will have full control over the database and this is the reason that they are moving to AWS in the first place "The company plans to quickly migrate to AWS because of limited resources for the database, backup administration, and data center maintenance?"

upvoted 1 times

pentium75 11 months, 1 week ago

"Amazon RDS Custom (B) is a managed database service for applications that require customization of the underlying operating system and database environment. Benefits of RDS automation with the access needed for legacy, packaged, and custom applications."

upvoted 2 times

Efren 1 year, 6 months ago

Selected Answer: B

RDS Custom when is something related to 3rd vendor, for me

upvoted 2 times

nosense 1 year, 6 months ago

not sure, but b probably

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #450

Topic 1


A company has a three-tier web application that is in a single server. The company wants to migrate the application to the AWS Cloud. The company also wants the application to align with the AWS Well-Architected Framework and to be consistent with AWS recommended best practices for security, scalability, and resiliency.

Which combination of solutions will meet these requirements? (Choose three.)

- A. Create a VPC across two Availability Zones with the application's existing architecture. Host the application with existing architecture on an Amazon EC2 instance in a private subnet in each Availability Zone with EC2 Auto Scaling groups. Secure the EC2 instance with security groups and network access control lists (network ACLs).
- B. Set up security groups and network access control lists (network ACLs) to control access to the database layer. Set up a single Amazon RDS database in a private subnet.
- C. Create a VPC across two Availability Zones. Refactor the application to host the web tier, application tier, and database tier. Host each tier on its own private subnet with Auto Scaling groups for the web tier and application tier. **Most Voted**
- D. Use a single Amazon RDS database. Allow database access only from the application tier security group.
- E. Use Elastic Load Balancers in front of the web tier. Control access by using security groups containing references to each layer's security groups. **Most Voted**
- F. Use an Amazon RDS database Multi-AZ cluster deployment in private subnets. Allow database access only from application tier security groups. **Most Voted**

Correct Answer: CEF

Community vote distribution



Comments

awsgeek75 **Highly Voted** 11 months ago

Selected Answer: CEF

The wording on this question makes things ambiguous for C. But, remember well-architected so:

A: Not ideal as it is suggesting using existing architecture but with autoscaling EC2. Doesn't leave room for improvement on scaling or reliability on each tier.

B: Single RDS, not well-architected

D: Again, single RDS

E,F are good options and C is only remaining good one.
upvoted 8 times

awsgeek75 11 months ago

C is badly worded IMHO because of this part " Refactor the application to host the web tier, application tier, and database tier." The database tier tier just makes it confusing when you don't read E and F.
upvoted 2 times

Abrar2022 Highly Voted 1 year, 6 months ago

Selected Answer: CEF

C-scalable and resilient
E-high availability of the application
F-Multi-AZ configuration provides high availability
upvoted 6 times

4919c76 Most Recent 1 month, 2 weeks ago

Selected Answer: EF

The web tier can only be in a public subnet so C is wrong
upvoted 1 times

ChhatwaniB 1 month, 3 weeks ago

in Option C , it saying all layer should be in it's own private subnet . Isn't the web layer should be in public subnet ?
upvoted 1 times

ChhatwaniB 1 month, 3 weeks ago

in Option C , it saying all layer should be in it's own private subnet . Isn't the web layer should be in public subnet ?
upvoted 1 times

Burrito69 8 months, 2 weeks ago

remove singles and remove network ACLs
upvoted 3 times

jjcode 10 months, 2 weeks ago

i would flag this on the test and do it last.
upvoted 5 times

argl1995 1 year, 5 months ago

option A cannot be the answer as Security group is at instance level whereas a NACL is at the subnet level. Having said that option C is the right one as the VPC cannot span across the regions and here it is mentioned two AZs for which I am guessing it is a default VPC which is created in each region with a subnet in each AZ.
upvoted 2 times

argl1995 1 year, 5 months ago

So, CEF is the right answer
upvoted 2 times

Gooniegoogoo 1 year, 5 months ago

How can you create a VPC across 2 AZ? i only see EF here.. if they mean 2 separate VPC then that is different but a VPC cannot span two AZ..
upvoted 1 times

lemur88 1 year, 3 months ago

A VPC most definitely can span across 2 AZ. You may be thinking of subnets.
upvoted 3 times

marufxplorer 1 year, 5 months ago

I also agree with CEF but chatGPT answer is ACE. A and C is the similar

Another Logic F is not True because in the question not mentioned about DB
upvoted 1 times

awsgeek75 10 months, 3 weeks ago

ChatGPT is a language parser. It is not an AWS solution architect!
upvoted 3 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: CEF

CEF is best
upvoted 2 times

antropaws 1 year, 6 months ago

Selected Answer: CEF

It's clearly CEF.
upvoted 2 times

omoakin 1 year, 6 months ago

B- to control access to database
C-scalable and resilient
E-high availability of the application
upvoted 1 times

lucdt4 1 year, 6 months ago

Selected Answer: CEF

CEF
A: application's existing architecture is wrong (single AZ)
B: single AZ
D: Single AZ
upvoted 2 times

cloudenthusiast 1 year, 6 months ago

C.
This solution follows the recommended architecture pattern of separating the web, application, and database tiers into different subnets. It provides better security, scalability, and fault tolerance.
E.By using Elastic Load Balancers (ELBs), you can distribute traffic to multiple instances of the web tier, increasing scalability and availability. Controlling access through security groups allows for fine-grained control and ensures only authorized traffic reaches each layer.
F.
Deploying an Amazon RDS database in a Multi-AZ configuration provides high availability and automatic failover. Placing the database in private subnets enhances security. Allowing database access only from the application tier security groups limits exposure and follows the principle of least privilege.
upvoted 5 times

mwwt2022 11 months ago

good explanation
upvoted 2 times

nosense 1 year, 6 months ago

Selected Answer: CEF

Only this valid for best practices and well architected
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #451

Topic 1


A company is migrating its applications and databases to the AWS Cloud. The company will use Amazon Elastic Container Service (Amazon ECS), AWS Direct Connect, and Amazon RDS.

Which activities will be managed by the company's operational team? (Choose three.)

- A. Management of the Amazon RDS infrastructure layer, operating system, and platforms
- B. Creation of an Amazon RDS DB instance and configuring the scheduled maintenance window **Most Voted**
- C. Configuration of additional software components on Amazon ECS for monitoring, patch management, log management, and host intrusion detection **Most Voted**
- D. Installation of patches for all minor and major database versions for Amazon RDS
- E. Ensure the physical security of the Amazon RDS infrastructure in the data center
- F. Encryption of the data that moves in transit through Direct Connect **Most Voted**

Correct Answer: BCF

Community vote distribution



Comments

pentium75 **Highly Voted** 11 months, 1 week ago

Selected Answer: BCF

ADE = AWS responsibility
upvoted 10 times

awsgeek75 **Highly Voted** 10 months, 3 weeks ago

Selected Answer: BCF

Just to clarify on F. Direct Connect is an ISP and AWS offering, I consider it as a physical connection just like you get from your ISP at home. There is not security on it until you build security on the connection. AWS provides Direct Connect but it does not provide encryption level security on data movement through it by default. It's the customer's responsibility.
upvoted 9 times

Guru4Cloud **Most Recent** 1 year, 3 months ago

Selected Answer: BCF

B: Creating an RDS instance and configuring the maintenance window is done by the customer.

C: Adding monitoring, logging, etc on ECS is managed by the customer.

F: Encrypting Direct Connect traffic is handled by the customer.

upvoted 5 times

james2033 1 year, 4 months ago

Selected Answer: BCF

In question has 3 keyword "Amazon ECS", "AWS Direct Connect", "Amazon RDS". With per Amazon services, choose 1 according answer. Has 6 items, need pick 3 items.

ECS --> choose C.

Direct Connect --> choose F.

RDS --> Exclude A (by keyword "infrastructure layer"), Choose B. Exclusive D (by keyword "patches for all minor and major database versions for Amazon RDS"). Exclusive E (by keyword "Ensure the physical security of the Amazon RDS"). Easy question.

upvoted 5 times

kapit 1 year, 5 months ago

BC & F (no automatic encryption with direct connect

upvoted 2 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: BF

Amazon ECS is a fully managed service, the ops team only focus on building their applications, not the environment.

Only option B and F makes sense.

upvoted 1 times

pentium75 11 months, 1 week ago

Plus C (we were asked for three). Configuration (!) of components for monitoring, log management etc.; those services exist from AWS but you need to configure them (which logs do you want to store for how long etc.).

upvoted 2 times

antropaws 1 year, 6 months ago

Selected Answer: BCF

100% BCF.

upvoted 2 times

lucdt4 1 year, 6 months ago

Selected Answer: BCF

BCF

B: Mentioned RDS

C: Mentioned ECS

F: Mentioned Direct connect

upvoted 5 times

hiroohiroo 1 year, 6 months ago

Selected Answer: BCF

Yes BCF

upvoted 2 times

omoakin 1 year, 6 months ago

I agree BCF

upvoted 2 times

nosense 1 year, 6 months ago

Selected Answer: BCF

Bcf for me
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #452

Topic 1

A company runs a Java-based job on an Amazon EC2 instance. The job runs every hour and takes 10 seconds to run. The job runs on a scheduled interval and consumes 1 GB of memory. The CPU utilization of the instance is low except for short surges during which the job uses the maximum CPU available. The company wants to optimize the costs to run the job.

Which solution will meet these requirements?

- A. Use AWS App2Container (A2C) to containerize the job. Run the job as an Amazon Elastic Container Service (Amazon ECS) task on AWS Fargate with 0.5 virtual CPU (vCPU) and 1 GB of memory.
- B. Copy the code into an AWS Lambda function that has 1 GB of memory. Create an Amazon EventBridge scheduled rule to run the code each hour. **Most Voted**
- C. Use AWS App2Container (A2C) to containerize the job. Install the container in the existing Amazon Machine Image (AMI). Ensure that the schedule stops the container when the task finishes.
- D. Configure the existing schedule to stop the EC2 instance at the completion of the job and restart the EC2 instance when the next job starts.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**omarshaban** **Highly Voted** 1 year, 4 months ago

THIS WAS IN MY EXAM

upvoted 9 times

awsgeek75 **Highly Voted** 1 year, 4 months ago**Selected Answer: B**

Never done it myself but apparently you can run Java in Lambda all the way to latest version
<https://docs.aws.amazon.com/lambda/latest/dg/lambda-java.html>

upvoted 6 times

Danilus **Most Recent** 7 months ago**Selected Answer: B**

key -The company wants to optimize the costs to run the job.
the answer is B because lambda provides the most cost-effective,easy to mange and scalable solution.it minimizes the need to manage infrastructure because you only pay for the compute time used during job execution

upvoted 2 times

noircesar25 1 year, 3 months ago

can someone explain what makes A wrong, im aware that C hasnt covered all the requirements but A seems good with fargate serverless and autoscaling functionalities, plus AWS app2container is for .NET and JAVA

upvoted 4 times

bujuman 1 year, 2 months ago

Statement: A company runs a Java-based job on an Amazon EC2 instance

Requirement: The company wants to optimize the costs to run the job

Regarding option A: App2Container is more likely for migratiing legacy application to conatiner based application. Which is not the main purpose of this use case.

We are asked to reduce cost on a application that is already running under EC2 instance.

So option B has the hight weight, cause lambda could perfectly do the job with the minimal cost

upvoted 4 times

JA2018 6 months, 2 weeks ago

you only need to pay for the compute resources consumed as and when each Lambda function is actually running (for the running duration)

upvoted 1 times

Murtadhaceit 1 year, 6 months ago

Selected Answer: B

This question is intended for Lambda. Just searched for Lambda with Event bridge. I

upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: B

Lambda allows you to allocate memory for your functions in increments of 1 MB, ranging from a minimum of 128 MB to a maximum of 10,240 MB (10 GB).

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

Remember - AWS Lambda function can go up to 10 GB of memory, instead of free tier only allow 512MB.

upvoted 6 times

james2033 1 year, 10 months ago

Selected Answer: B

"AWS Batch jobs as EventBridge targets" at <https://docs.aws.amazon.com/batch/latest/userguide/batch-cwe-target.html>

AWS Batch + Amazon EventBridge <https://docs.aws.amazon.com/batch/latest/userguide/batch-cwe-target.html> .

AWS Lambda just for a point of time per period. Choose B.

upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: B

10 seconds to run, optimize the costs, consumes 1 GB of memory = AWS Lambda function.

upvoted 3 times

alexandercamachop 2 years ago

Selected Answer: B

AWS Lambda automatically scales resources to handle the workload, so you don't have to worry about managing the underlying infrastructure. It provisions the necessary compute resources based on the configured memory size (1 GB in this

case) and executes the job in a serverless environment.

By using Amazon EventBridge, you can create a scheduled rule to trigger the Lambda function every hour, ensuring that the job runs on the desired interval.

upvoted 2 times

Yadav_Sanjay 2 years ago

Selected Answer: B

B - Within 10 sec and 1 GB Memory (Lambda Memory 128MB to 10GB)

upvoted 3 times

Yadav_Sanjay 2 years ago

<https://docs.aws.amazon.com/lambda/latest/operatorguide/computing-power.html>

upvoted 2 times

Efren 2 years ago

Selected Answer: B

Agreed, B Lambda

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #453

Topic 1

A company wants to implement a backup strategy for Amazon EC2 data and multiple Amazon S3 buckets. Because of regulatory requirements, the company must retain backup files for a specific time period. The company must not alter the files for the duration of the retention period.

Which solution will meet these requirements?

- A. Use AWS Backup to create a backup vault that has a vault lock in governance mode. Create the required backup plan.
- B. Use Amazon Data Lifecycle Manager to create the required automated snapshot policy.
- C. Use Amazon S3 File Gateway to create the backup. Configure the appropriate S3 Lifecycle management.
- D. Use AWS Backup to create a backup vault that has a vault lock in compliance mode. Create the required backup plan.

Most Voted

Correct Answer: D

Community vote distribution

D (100%)

Comments

Efren Highly Voted 2 years ago

D, Governance is like the government, they can do things you cannot, like delete files or backups :D Compliance, nobody can!
upvoted 39 times

cmbt 1 year, 11 months ago

Finally I understood!
upvoted 4 times

joshnort 1 year, 11 months ago

Great analogy
upvoted 8 times

nosense Highly Voted 2 years ago

Selected Answer: D

D bcs in governance we can delete backup

upvoted 5 times

Danilus Most Recent 7 months ago

Selected Answer: D

key-The company must not alter the files for the duration of the retention period

the answer is D because backup with vault lock in compliance mode make sure that the files when are locked they cannot be deleted or modified for the specified retention period it makes the files immutable

why is not A? because the governance mode... with this mode backups allows flexibility to alter or delete backups

upvoted 2 times

f2e2419 1 year, 5 months ago

Selected Answer: D

D. Use AWS Backup to create a backup vault that has a vault lock in compliance mode. Create the required backup plan

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

D. Use AWS Backup to create a backup vault that has a vault lock in compliance mode. Create the required backup plan

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

Use AWS Backup to create a backup vault that has a vault lock in compliance mode. Create the required backup plan

upvoted 3 times

ccat91 1 year, 10 months ago

Selected Answer: D

Compliance mode

upvoted 2 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: D

Must not alter the files for the duration of the retention period = Compliance Mode

upvoted 2 times

antropaws 2 years ago

Selected Answer: D

D for sure.

upvoted 2 times

dydzah 2 years ago

Selected Answer: D

<https://docs.aws.amazon.com/aws-backup/latest/devguide/vault-lock.html>

upvoted 3 times

cloudenthusiast 2 years ago

Selected Answer: D

compliance mode

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #454

Topic 1


A company has resources across multiple AWS Regions and accounts. A newly hired solutions architect discovers a previous employee did not provide details about the resources inventory. The solutions architect needs to build and map the relationship details of the various workloads across all accounts.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. Use AWS Systems Manager Inventory to generate a map view from the detailed view report.
- B. Use AWS Step Functions to collect workload details. Build architecture diagrams of the workloads manually.
- C. Use Workload Discovery on AWS to generate architecture diagrams of the workloads. **Most Voted**
- D. Use AWS X-Ray to view the workload details. Build architecture diagrams with relationships.

Correct Answer: C

Community vote distribution



Comments

Danilus 7 months ago

Key - Build and map the relationship details of the various workloads.
Key - MOST operationally efficient way.

The answer is C because Workload Discovery is a tool that allows you to find automatically the workloads on the cloud, giving you a visual representation of the infrastructure and the relationship between the resources through different accounts.

upvoted 2 times

zinabu 1 year, 2 months ago

workload discovery=architecture diagram

upvoted 3 times

osmk 1 year, 4 months ago

<https://docs.aws.amazon.com/solutions/latest/workload-discovery-on-aws/solution-overview.html> Workload Discovery on AWS is a visualization tool that automatically generates architecture diagrams of your workload on AWS. You can use this solution to build, customize, and share detailed workload visualizations based on live data from AWS

upvoted 3 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

A: Systems Manager Inventory -> Metadata
B: Not possible (correct me if I'm wrong)
D: X-Ray is for application debugging
C: Workload Discovery is purpose built tool for this type of usage
upvoted 4 times

NayeraB 1 year, 3 months ago

Even if B is possible, it has "manually" in it which we won't do because we're lazy in this question
upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: C

Workload Discovery on AWS (formerly called AWS Perspective) is a tool to visualize AWS Cloud workloads. Use Workload Discovery on AWS to build, customize, and share detailed architecture diagrams of your workloads based on live data from AWS.
upvoted 3 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: C

use Workload Discovery on AWS
upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: C

Workload Discovery is purpose-built to automatically generate visual mappings of architectures across accounts and Regions. This makes it the most operationally efficient way to meet the requirements.
upvoted 4 times

MrAWSAssociate 1 year, 11 months ago

Selected Answer: C

Option A: AWS SSM offers "Software inventory": Collect software catalog and configuration for your instances.
Option C: Workload Discovery on AWS: is a tool for maintaining an inventory of the AWS resources across your accounts and various Regions and mapping relationships between them, and displaying them in a web UI.
upvoted 4 times

DrWatson 2 years ago

Selected Answer: A

<https://aws.amazon.com/blogs/mt/visualizing-resources-with-workload-discovery-on-aws/>
upvoted 1 times

pentium75 1 year, 5 months ago

That is C
upvoted 2 times

Abrar2022 2 years ago

Selected Answer: C

AWS Workload Discovery - create diagram, map and visualise AWS resources across AWS accounts and Regions
upvoted 3 times

Abrar2022 2 years ago

Workload Discovery on AWS can map AWS resources across AWS accounts and Regions and visualize them in a UI provided on the website.
upvoted 2 times

hiroohiroo 2 years ago

Selected Answer: C

https://aws.amazon.com/jp/builders-flash/202209/workload-discovery-on-aws/?awsf.filter-name=*all

upvoted 3 times

omoakin 2 years ago

Only C makes sense

upvoted 3 times

cloudenthusiast 2 years ago

Selected Answer: C

Workload Discovery on AWS is a service that helps visualize and understand the architecture of your workloads across multiple AWS accounts and Regions. It automatically discovers and maps the relationships between resources, providing an accurate representation of the architecture.

upvoted 3 times

Efren 2 years ago

Not sure here tbh

To efficiently build and map the relationship details of various workloads across multiple AWS Regions and accounts, you can use the AWS Systems Manager Inventory feature in combination with AWS Resource Groups. Here's a solution that can help you achieve this:

AWS Systems Manager Inventory:

upvoted 1 times

nosense 2 years ago

Selected Answer: C

only c mapping relationships

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #455

Topic 1

A company uses AWS Organizations. The company wants to operate some of its AWS accounts with different budgets. The company wants to receive alerts and automatically prevent provisioning of additional resources on AWS accounts when the allocated budget threshold is met during a specific period.

Which combination of solutions will meet these requirements? (Choose three.)

A. Use AWS Budgets to create a budget. Set the budget amount under the Cost and Usage Reports section of the required AWS accounts.


B. Use AWS Budgets to create a budget. Set the budget amount under the Billing dashboards of the required AWS accounts. **Most Voted**

C. Create an IAM user for AWS Budgets to run budget actions with the required permissions.

D. Create an IAM role for AWS Budgets to run budget actions with the required permissions. **Most Voted**

E. Add an alert to notify the company when each account meets its budget threshold. Add a budget action that selects the IAM identity created with the appropriate config rule to prevent provisioning of additional resources.

F. Add an alert to notify the company when each account meets its budget threshold. Add a budget action that selects the IAM identity created with the appropriate service control policy (SCP) to prevent provisioning of additional resources.

Most Voted**Correct Answer: BDF***Community vote distribution***Comments****vesen22** **Highly Voted** 2 years ago**Selected Answer: BDF**

I don't see why adf has the most voted when almost everyone has chosen bdf, smh
https://acloudguru.com/videos/acg-fundamentals/how-to-set-up-an-aws-billing-and-budget-alert?utm_source=google&utm_medium=paid-search&utm_campaign=cloud-transformation&utm_term=ssi-global-acg-core-dsa&utm_content=free-trial&gclid=Cj0KCQjwmtGjBhDhARIsAEqfDEcDfXdLul2NxgSMxKracIITZimWotDBRpsJPpx8IS9T4NndKhbUqPlaAlzhEALw_wcB
upvoted 14 times

omarshaban Highly Voted 1 year, 4 months ago

IN MY EXAM

upvoted 9 times

FlyingHawk Most Recent 6 months ago

Selected Answer: B

AWS budget is under the Budget and Planning section in Billing and Cost Management console

upvoted 1 times

FlyingHawk 6 months ago

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/view-billing-dashboard.html>

upvoted 1 times

TariqKipkemei 1 year, 7 months ago

Selected Answer: DF

Its 11/Nov/2023. Options D&F are definitely required.

As for the budget, right from the aws console, the only place to set this up is:

AWS Billing>Cost Management>Budgets.

upvoted 5 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: BDF

How to create a budget:

Billing console > budget > create budget!

upvoted 6 times

Chris22usa 1 year, 11 months ago

ACF:

Option B is incorrect because the budget amount should be set under the Cost and Usage Reports section, not the Billing dashboards.

upvoted 1 times

pentium75 1 year, 5 months ago

"Create an AWS Budget: Go to the AWS Billing Dashboard"

<https://awslabs.github.io/scale-out-computing-on-aws/workshops/TKO-Scale-Out-Computing/modules/071-budgets/>

upvoted 1 times

Abrar2022 2 years ago

Selected Answer: BDF

How to create a budget:

Billing console > budget > create budget!

upvoted 3 times

udo2020 2 years ago

It is BDF because there is actually a Billing Dashboard available.

upvoted 8 times

hiroohiroo 2 years ago

Selected Answer: BDF

https://docs.aws.amazon.com/ja_jp/awsaccountbilling/latest/aboutv2/view-billing-dashboard.html

upvoted 6 times

y0 2 years ago

BDF - Budgets can be set from the billing dashboard in AWS console

upvoted 3 times

cloudenthusiast 2 years ago

Selected Answer: ADF

Currently, AWS does not have a specific feature called "AWS Billing Dashboards."

upvoted 6 times

[Removed] 2 years ago

<https://awslabs.github.io/scale-out-computing-on-aws/workshops/TKO-Scale-Out-Computing/modules/071-budgets/>

upvoted 4 times

Efren 2 years ago

if im not wrong, those are correct

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #456

Topic 1

A company runs applications on Amazon EC2 instances in one AWS Region. The company wants to back up the EC2 instances to a second Region. The company also wants to provision EC2 resources in the second Region and manage the EC2 instances centrally from one AWS account.

Which solution will meet these requirements MOST cost-effectively?

- A. Create a disaster recovery (DR) plan that has a similar number of EC2 instances in the second Region. Configure data replication.
- B. Create point-in-time Amazon Elastic Block Store (Amazon EBS) snapshots of the EC2 instances. Copy the snapshots to the second Region periodically.
- C. Create a backup plan by using AWS Backup. Configure cross-Region backup to the second Region for the EC2 instances.
Most Voted
- D. Deploy a similar number of EC2 instances in the second Region. Use AWS DataSync to transfer the data from the source Region to the second Region.

Correct Answer: C*Community vote distribution***Comments****cloudenthusiast** **Highly Voted** 2 years ago**Selected Answer: C**

Using AWS Backup, you can create backup plans that automate the backup process for your EC2 instances. By configuring cross-Region backup, you can ensure that backups are replicated to the second Region, providing a disaster recovery capability. This solution is cost-effective as it leverages AWS Backup's built-in features and eliminates the need for manual snapshot management or deploying and managing additional EC2 instances in the second Region.

upvoted 8 times

Guru4Cloud **Highly Voted** 1 year, 9 months ago**Selected Answer: C**

C is the most cost-effective solution that meets all the requirements.

AWS Backup provides automated backups across Regions for EC2 instances. This handles the backup requirement.

AWS Backup is more cost-effective for cross-Region EC2 backups than using EBS snapshots manually or DataSync.

upvoted 5 times

wizcloudifa 1 year, 1 month ago

Well Datasync is to transfer data from on-prem to AWS rather than AWS R1 to R2, and it is use for a one-time migration rather than repetitive data migration activities like backups hence its out(on so many levels)

upvoted 2 times

elmyth Most Recent 7 months ago

Selected Answer: B

We have implemented solution similar to B, it's really small and simple script, so we don't have much manual work here, and it's really the most cost-effective.

upvoted 3 times

Certibanksman 9 months, 2 weeks ago

Selected Answer: B

Option B (EBS snapshots with cross-Region copy) is the most cost-effective solution for backing up EC2 instances to a second Region while allowing for centralized management and easy recovery when needed.

upvoted 3 times

Daminaij 9 months, 3 weeks ago

The answer is c

upvoted 1 times

bogobob 1 year, 6 months ago

Selected Answer: D

How does AWS Backup address that "The company also wants to provision EC2 resources in the second Region"?

upvoted 3 times

NSA_Poker 1 year ago

With AWS Backup you can 'back up the EC2 instances to a second Region' by implementing cross-Region backup & 'provision EC2 resources in the second Region' by restoring the backup using the AWS Backup console.

<https://docs.aws.amazon.com/aws-backup/latest/devguide/restore-resource.html>

upvoted 1 times

pentium75 1 year, 5 months ago

How do A, B or D address that? They want to "provision EC2 resources", nobody says that this should be copies of the existing servers. And if it should be copies of the existing servers, wouldn't we need the same (not "a similar") number of servers? We have no idea how many applications on how many servers they have.

upvoted 3 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: C

AWS backup

upvoted 2 times

omoakin 2 years ago

CCCCC

. Create a backup plan by using AWS Backup. Configure cross-Region backup to the second Region for the EC2 instances.

upvoted 1 times

Blingy 2 years ago

CCCCCCC

upvoted 2 times

Efren 2 years ago

C, i would say same, always AWS Backup
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #457

Topic 1


A company that uses AWS is building an application to transfer data to a product manufacturer. The company has its own identity provider (IdP). The company wants the IdP to authenticate application users while the users use the application to transfer data. The company must use Applicability Statement 2 (AS2) protocol.

Which solution will meet these requirements?

- A. Use AWS DataSync to transfer the data. Create an AWS Lambda function for IdP authentication.
- B. Use Amazon AppFlow flows to transfer the data. Create an Amazon Elastic Container Service (Amazon ECS) task for IdP authentication.
- C. Use AWS Transfer Family to transfer the data. Create an AWS Lambda function for IdP authentication. **Most Voted**
- D. Use AWS Storage Gateway to transfer the data. Create an Amazon Cognito identity pool for IdP authentication.

Correct Answer: C

Community vote distribution



Comments

TariqKipkemei **Highly Voted** 1 year, 11 months ago

Selected Answer: C

Option C stands out stronger because AWS Transfer Family securely scales your recurring business-to-business file transfers to AWS Storage services using SFTP, FTPS, FTP, and AS2 protocols.

And AWS Lambda can be used to authenticate users with the company's IdP.

upvoted 10 times

baba365 1 year, 11 months ago

Ans : C

To authenticate your users, you can use your existing identity provider with AWS Transfer Family. You integrate your identity provider using an AWS Lambda function, which authenticates and authorizes your users for access to Amazon S3 or Amazon Elastic File System (Amazon EFS).

<https://docs.aws.amazon.com/transfer/latest/userguide/custom-identity-provider-users.html>

upvoted 6 times

Danilus Most Recent 7 months ago

Selected Answer: C

key-The company must use Applicability Statement 2 (AS2) protocol

The only service in AWS that accepts this protocol is AWS Transfer Family, which provides a managed solution for secure file transfers using AS2, along with essential features such as encryption, digital signatures, and message disposition notifications (MDN)

upvoted 2 times

zinabu 1 year, 2 months ago

aws transfer family for data transfer and lamda function for idp authentication

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: C

<https://aws.amazon.com/about-aws/whats-new/2022/07/aws-transfer-family-support-applicability-statement-2-as2/>

upvoted 4 times

potomac 1 year, 7 months ago

Selected Answer: C

To authenticate your users, you can use your existing identity provider with AWS Transfer Family. You integrate your identity provider using an AWS Lambda function, which authenticates and authorizes your users for access to Amazon S3 or Amazon Elastic File System (Amazon EFS).

upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: C

Applicability Statement 2 (AS2) is a business-to-business (B2B) messaging protocol used to exchange Electronic Data Interchange (EDI) documents. With AWS Transfer Family's AS2 capabilities, you can securely exchange AS2 messages at scale while maintaining compliance and interoperability with your trading partners.

upvoted 4 times

thanhnv142 1 year, 7 months ago

D is ok

upvoted 1 times

hsinchang 1 year, 10 months ago

its own IdP -> Lambda

upvoted 3 times

dydzah 2 years ago

Selected Answer: C

<https://docs.aws.amazon.com/transfer/latest/userguide/custom-identity-provider-users.html>

upvoted 2 times

examtopictempacc 2 years ago

Selected Answer: C

C is correct. AWS Transfer Family supports the AS2 protocol, which is required by the company. Also, AWS Lambda can be used to authenticate users with the company's IdP, which meets the company's requirement.

upvoted 3 times

EA100 2 years ago

Answer - D

AS2 is a widely used protocol for secure and reliable data transfer. In this scenario, the company wants to transfer data using the AS2 protocol and authenticate application users using their own identity provider (IdP). AWS Storage Gateway provides a hybrid cloud storage solution that enables data transfer between on-premises environments and AWS.

By using AWS Storage Gateway, you can set up a gateway that supports the AS2 protocol for data transfer. Additionally, you can configure authentication using an Amazon Cognito identity pool. Amazon Cognito provides a comprehensive authentication and user management service that integrates with various identity providers, including your own IdP.

authentication and user management service that integrates with various identity providers, including your own IdP.

Therefore, Option D is the correct solution as it leverages AWS Storage Gateway for AS2 data transfer and allows authentication using an Amazon Cognito identity pool integrated with the company's IdP.

upvoted 2 times

deechean 1 year, 9 months ago

AWS Transfer Family also support AS2

upvoted 2 times

hiroohiroo 2 years ago

Selected Answer: C

<https://repost.aws/articles/ARo2ihKKThT2Cue5j6yVUgsQ/articles/ARo2ihKKThT2Cue5j6yVUgsQ/aws-transfer-family-announces-support-for-sending-as2-messages-over-https?>

upvoted 1 times

omoakin 2 years ago

C is correct

upvoted 1 times

omoakin 2 years ago

This is a new Qtn n AS2 is newly supported by AWS Transfer family.....good timing to know ur stuffs.

upvoted 2 times

nosense 2 years ago

Option D looks the better option because it is more secure, scalable, cost-effective, and easy to use than option C.

upvoted 1 times

cloudenthusiast 2 years ago

Selected Answer: D

AWS Storage Gateway supports the AS2 protocol for transferring data. By using AWS Storage Gateway, the company can integrate its own IdP authentication by creating an Amazon Cognito identity pool. Amazon Cognito provides user authentication and authorization capabilities, allowing the company to authenticate application users using its own IdP.

AWS Transfer Family does not currently support the AS2 protocol. AS2 is a specific protocol used for secure and reliable data transfer, often used in business-to-business (B2B) scenarios. In this case, option C, which suggests using AWS Transfer Family, would not meet the requirement of using the AS2 protocol.

upvoted 3 times

omoakin 2 years ago

AWS Transfer Family now supports the Applicability Statement 2 (AS2) protocol, complementing existing protocol support for SFTP, FTPS, and FTP

upvoted 2 times

y0 2 years ago

This is not a case for storage gateway which is more used for a hybrid like environment. Here, to transfer data, we can think of Datasync or Transfer family and considering AS2 protocol, transfer family looks good

upvoted 3 times

Efren 2 years ago

ChatGP

To meet the requirements of using an identity provider (IdP) for user authentication and the AS2 protocol for data transfer, you can implement the following solution:

AWS Transfer Family: Use AWS Transfer Family, specifically AWS Transfer for SFTP or FTPS, to handle the data transfer using the AS2 protocol. AWS Transfer for SFTP and FTPS provide fully managed, highly available SFTP and FTPS servers in the AWS Cloud.

Not sure about Lambda tho

upvoted 3 times

Efren 2 years ago

Maybe yes

The Lambda authorizer authenticates the token with the third-party identity provider.
upvoted 2 times

cloudenthusiast 2 years ago

Also from ChatGPT

AWS Transfer Family supports multiple protocols, including AS2, and can be used for data transfer. By utilizing AWS Transfer Family, the company can integrate its own IdP authentication by creating an AWS Lambda function.

Both options D and C are valid solutions for the given requirements. The choice between them would depend on additional factors such as specific preferences, existing infrastructure, and overall architectural considerations.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #458

Topic 1

A solutions architect is designing a RESTAPI in Amazon API Gateway for a cash payback service. The application requires 1 GB of memory and 2 GB of storage for its computation resources. The application will require that the data is in a relational format.

Which additional combination of AWS services will meet these requirements with the LEAST administrative effort? (Choose two.)

A. Amazon EC2

B. AWS Lambda **Most Voted**

C. Amazon RDS **Most Voted**

D. Amazon DynamoDB

E. Amazon Elastic Kubernetes Services (Amazon EKS)

Correct Answer: BC

Community vote distribution

BC (89%)

AC (11%)

Comments

cloudenthusiast **Highly Voted** 2 years ago

Selected Answer: BC

"The application will require that the data is in a relational format" so DynamoDB is out. RDS is the choice. Lambda is serverless.
upvoted 17 times

smurfing2k17 **Most Recent** 1 year, 1 month ago

Why cant it be AC? We don't know the time of job runs right?
upvoted 2 times

JA2018 6 months, 2 weeks ago

requirement/ key not stated in stem
upvoted 1 times

TarinKinkamei 1 year, 11 months ago

tanquaphemer 1 year, 11 months ago

Selected Answer: BC

AWS Lambda and Amazon RDS

upvoted 3 times

handsonlabsaws 2 years ago

Selected Answer: AC

"2 GB of storage for its COMPUTATION resources" the maximum for Lambda is 512MB.

upvoted 3 times

PLN6302 1 year, 9 months ago

Lambda now supports upto 10GB of memory

upvoted 8 times

Kp88 1 year, 10 months ago

I thought the same but seems like you can go all the way to 10gb. 512mb is the free tier

<https://docs.aws.amazon.com/lambda/latest/dg/configuration-function-common.html#configuration-ephemeral-storage>

upvoted 3 times

r3mo 1 year, 12 months ago

At first I was thinking the same. But the computation memory for the lambda function is 1gb not 2gb. Hence, if you go to basic settings when you create the lambda function you can select a in the memory settings the 1024 MB (1Gb) and that solve the problem.

upvoted 2 times

Efren 2 years ago

Selected Answer: BC

Relational Data RDS and computing for Lambda

upvoted 4 times

nosense 2 years ago

bc for me

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #459

Topic 1

A company uses AWS Organizations to run workloads within multiple AWS accounts. A tagging policy adds department tags to AWS resources when the company creates tags.

An accounting team needs to determine spending on Amazon EC2 consumption. The accounting team must determine which departments are responsible for the costs regardless of AWS account. The accounting team has access to AWS Cost Explorer for all AWS accounts within the organization and needs to access all reports from Cost Explorer.

Which solution meets these requirements in the MOST operationally efficient way?

- A. From the Organizations management account billing console, activate a user-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by tag name, and filter by EC2. **Most Voted**
- B. From the Organizations management account billing console, activate an AWS-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by tag name, and filter by EC2.
- C. From the Organizations member account billing console, activate a user-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by the tag name, and filter by EC2.
- D. From the Organizations member account billing console, activate an AWS-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by tag name, and filter by EC2.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**cloudenthusiast** **Highly Voted** 1 year, 6 months ago**Selected Answer: A**

By activating a user-defined cost allocation tag named "department" and creating a cost report in Cost Explorer that groups by the tag name and filters by EC2, the accounting team will be able to track and attribute costs to specific departments across all AWS accounts within the organization. This approach allows for consistent cost allocation and reporting regardless of the AWS account structure.

upvoted 9 times

luisgu **Highly Voted** 1 year, 6 months ago

Selected Answer: A

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/custom-tags.html>

upvoted 6 times

awsgeek75 **Most Recent** 11 months ago

Selected Answer: A

Management not user.

upvoted 5 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

From the Organizations management account billing console, activate a user-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by tag name, and filter by EC2.

upvoted 5 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: A

From the Organizations management account billing console, activate a user-defined cost allocation tag named department. Create one cost report in Cost Explorer grouping by tag name, and filter by EC2.

upvoted 3 times

hiroohiroo 1 year, 6 months ago

Selected Answer: A

https://docs.aws.amazon.com/ja_jp/awsaccountbilling/latest/aboutv2/activating-tags.html

upvoted 4 times

nosense 1 year, 6 months ago

Selected Answer: A

a for me

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #460

Topic 1

A company wants to securely exchange data between its software as a service (SaaS) application Salesforce account and Amazon S3. The company must encrypt the data at rest by using AWS Key Management Service (AWS KMS) customer managed keys (CMKs). The company must also encrypt the data in transit. The company has enabled API access for the Salesforce account.

- A. Create AWS Lambda functions to transfer the data securely from Salesforce to Amazon S3.
- B. Create an AWS Step Functions workflow. Define the task to transfer the data securely from Salesforce to Amazon S3.
- C. Create Amazon AppFlow flows to transfer the data securely from Salesforce to Amazon S3. **Most Voted**
- D. Create a custom connector for Salesforce to transfer the data securely from Salesforce to Amazon S3.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**cloudenthusiast** **Highly Voted** 2 years ago**Selected Answer: C**

Amazon AppFlow is a fully managed integration service that allows you to securely transfer data between different SaaS applications and AWS services. It provides built-in encryption options and supports encryption in transit using SSL/TLS protocols. With AppFlow, you can configure the data transfer flow from Salesforce to Amazon S3, ensuring data encryption at rest by utilizing AWS KMS CMKs.

upvoted 15 times

Guru4Cloud **Highly Voted** 1 year, 9 months ago**Selected Answer: C**

- Amazon AppFlow can securely transfer data between Salesforce and Amazon S3.
- AppFlow supports encrypting data at rest in S3 using KMS CMKs.
- AppFlow supports encrypting data in transit using HTTPS/TLS.
- AppFlow provides built-in support and templates for Salesforce and S3, requiring less custom configuration than solutions like Lambda, Step Functions, or custom connectors.
- So Amazon AppFlow is the easiest way to meet all the requirements of securely transferring data between Salesforce and S3 with encryption at rest and in transit.

upvoted 9 times

1a22522 **Most Recent** 10 months 1 week ago

tezzj22 10 months, 1 week ago

Selected Answer: C

i do like myself some appflow flow
upvoted 2 times

zinabu 1 year, 2 months ago

SAAS=aws appflow
upvoted 4 times

cvoiceip 1 year, 5 months ago

Ans : C
Salesforce -----> Amazon AppFlow -----> S3
upvoted 3 times

hsinchang 1 year, 10 months ago

securely transfer data between Software-as-a-Service (SaaS) applications and AWS -> AppFlow
upvoted 3 times

TariqKipkemei 1 year, 11 months ago

Selected Answer: C

With Amazon AppFlow automate bi-directional data flows between SaaS applications and AWS services in just a few clicks
upvoted 2 times

DrWatson 2 years ago

Selected Answer: C

<https://docs.aws.amazon.com/appflow/latest/userguide/what-is-appflow.html>
upvoted 3 times

Abrar2022 2 years ago

All you need to know is that AWS AppFlow securely transfers data between different SaaS applications and AWS services
upvoted 3 times

hiroohiroo 2 years ago

Selected Answer: C

<https://docs.aws.amazon.com/appflow/latest/userguide/salesforce.html>
upvoted 4 times

Efren 2 years ago

Selected Answer: C

SaaS with another service, AppFlow
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #461

Topic 1

A company is developing a mobile gaming app in a single AWS Region. The app runs on multiple Amazon EC2 instances in an Auto Scaling group. The company stores the app data in Amazon DynamoDB. The app communicates by using TCP traffic and UDP traffic between the users and the servers. The application will be used globally. The company wants to ensure the lowest possible latency for all users.

Which solution will meet these requirements?

A. Use AWS Global Accelerator to create an accelerator. Create an Application Load Balancer (ALB) behind an accelerator endpoint that uses Global Accelerator integration and listening on the TCP and UDP ports. Update the Auto Scaling group to register instances on the ALB.

B. Use AWS Global Accelerator to create an accelerator. Create a Network Load Balancer (NLB) behind an accelerator endpoint that uses Global Accelerator integration and listening on the TCP and UDP ports. Update the Auto Scaling group to register instances on the NLB. **Most Voted**

C. Create an Amazon CloudFront content delivery network (CDN) endpoint. Create a Network Load Balancer (NLB) behind the endpoint and listening on the TCP and UDP ports. Update the Auto Scaling group to register instances on the NLB. Update CloudFront to use the NLB as the origin.

D. Create an Amazon CloudFront content delivery network (CDN) endpoint. Create an Application Load Balancer (ALB) behind the endpoint and listening on the TCP and UDP ports. Update the Auto Scaling group to register instances on the ALB. Update CloudFront to use the ALB as the origin.

Correct Answer: B

Community vote distribution

B (100%)

Comments

zinabu **Highly Voted** 8 months ago

TCP/UDP/IP based communication with server =NLB
for global low latency communication if IP/udp/tcp based = aws global accelerator
upvoted 5 times

sandordini **Most Recent** 7 months, 3 weeks ago

Selected Answer: B

Mobile gaming, UDP > AWS Global Accelerator + NLB
upvoted 3 times

Mikado211 11 months, 3 weeks ago

Selected Answer: B

UDP == NLB
NLB can't be used with Cloudfront, so we have to play with AWS Global accelerator
upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Use AWS Global Accelerator to create an accelerator. Create a Network Load Balancer (NLB) behind an accelerator endpoint that uses Global Accelerator integration and listening on the TCP and UDP ports. Update the Auto Scaling group to register instances on the NLB
upvoted 3 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: B

TCP and UDP = global accelerator and Network Load Balancer
upvoted 3 times

antropaws 1 year, 6 months ago

Selected Answer: B

Clearly B.
upvoted 2 times

eddie5049 1 year, 6 months ago

Selected Answer: B

NLB + Accelerator
upvoted 4 times

hiroohiroo 1 year, 6 months ago

Selected Answer: B

AWS Global Accelerator+NLB
upvoted 4 times

Efren 1 year, 6 months ago

Selected Answer: B

UDP, Global Accelerator plus NLB
upvoted 2 times

nosense 1 year, 6 months ago

Selected Answer: B

AWS Global Accelerator is a better solution for the mobile gaming app than CloudFront
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #462

Topic 1

A company has an application that processes customer orders. The company hosts the application on an Amazon EC2 instance that saves the orders to an Amazon Aurora database. Occasionally when traffic is high the workload does not process orders fast enough.

What should a solutions architect do to write the orders reliably to the database as quickly as possible?

- A. Increase the instance size of the EC2 instance when traffic is high. Write orders to Amazon Simple Notification Service (Amazon SNS). Subscribe the database endpoint to the SNS topic.
- B. Write orders to an Amazon Simple Queue Service (Amazon SQS) queue. Use EC2 instances in an Auto Scaling group behind an Application Load Balancer to read from the SQS queue and process orders into the database. **Most Voted**
- C. Write orders to Amazon Simple Notification Service (Amazon SNS). Subscribe the database endpoint to the SNS topic. Use EC2 instances in an Auto Scaling group behind an Application Load Balancer to read from the SNS topic.
- D. Write orders to an Amazon Simple Queue Service (Amazon SQS) queue when the EC2 instance reaches CPU threshold limits. Use scheduled scaling of EC2 instances in an Auto Scaling group behind an Application Load Balancer to read from the SQS queue and process orders into the database.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**cloudenthusiast** **Highly Voted** 1 year, 6 months ago**Selected Answer: B**

By decoupling the write operation from the processing operation using SQS, you ensure that the orders are reliably stored in the queue, regardless of the processing capacity of the EC2 instances. This allows the processing to be performed at a scalable rate based on the available EC2 instances, improving the overall reliability and speed of order processing.

upvoted 12 times

omarshaban **Highly Voted** 10 months, 3 weeks ago

IN MY EXAM

upvoted 9 times

Exam 401 **Most Recent** 1 year, 2 months ago

Guru4Cloud most recent 1 year, 3 months ago

Selected Answer: B

Decoupling the order processing from the application using Amazon SQS and leveraging Auto Scaling to handle the processing of orders based on the workload in the SQS queue is indeed the most efficient and scalable approach. This architecture addresses both reliability and performance concerns during traffic spikes.

upvoted 4 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: B

Write orders to an Amazon Simple Queue Service (Amazon SQS) queue. Use EC2 instances in an Auto Scaling group behind an Application Load Balancer to read from the SQS queue and process orders into the database.

upvoted 3 times

antropaws 1 year, 6 months ago

Selected Answer: B

100% B.

upvoted 3 times

omoakin 1 year, 6 months ago

BBBBBBBBBB

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #463

Topic 1

An IoT company is releasing a mattress that has sensors to collect data about a user's sleep. The sensors will send data to an Amazon S3 bucket. The sensors collect approximately 2 MB of data every night for each mattress. The company must process and summarize the data for each mattress. The results need to be available as soon as possible. Data processing will require 1 GB of memory and will finish within 30 seconds.

Which solution will meet these requirements MOST cost-effectively?

- A. Use AWS Glue with a Scala job
- B. Use Amazon EMR with an Apache Spark script
- C. Use AWS Lambda with a Python script **Most Voted**
- D. Use AWS Glue with a PySpark job

Correct Answer: C*Community vote distribution*

C (100%)

Comments**cloudenthusiast** **Highly Voted** 1 year, 6 months ago**Selected Answer: C**

AWS Lambda charges you based on the number of invocations and the execution time of your function. Since the data processing job is relatively small (2 MB of data), Lambda is a cost-effective choice. You only pay for the actual usage without the need to provision and maintain infrastructure.

upvoted 7 times

joechen2023 1 year, 5 months ago

but the question states "Data processing will require 1 GB of memory and will finish within 30 seconds." so it can't be C as Lambda support maximum 512M

upvoted 2 times

BillaRanga 10 months ago

Lambda can support upto 10 GB, But 512M is under free tier

upvoted 3 times

nilandd44gg 1 year, 4 months ago

C is valid.

Lambda quotas:

Memory - 128 MB to 10,240 MB, in 1-MB increments.

Note: Lambda allocates CPU power in proportion to the amount of memory configured. You can increase or decrease the memory and CPU power allocated to your function using the Memory (MB) setting. At 1,769 MB, a function has the equivalent of one vCPU.

Function timeout 900 seconds (15 minutes)

4 KB, for all environment variables associated with the function, in aggregate

<https://docs.aws.amazon.com/lambda/latest/dg/gettingstarted-limits.html>

upvoted 3 times

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: C

The data processing is lightweight, only requiring 1 GB memory and finishing in under 30 seconds. Lambda is designed for short, transient workloads like this.

Lambda scales automatically, invoking the function as needed when new data arrives. No servers to manage.

Lambda has a very low cost. You only pay for the compute time used to run the function, billed in 100ms increments. Much cheaper than provisioning EMR or Glue.

Processing can begin as soon as new data hits the S3 bucket by triggering the Lambda function. Provides low latency.

upvoted 5 times

Chiquitabandita **Most Recent** 1 year, 1 month ago

I understand C is a common answer "throw Lambda" seems to be a common theme for questions that need processing under 15 minutes for the test. but in reality, can the other solutions be viable options as well?

upvoted 4 times

Mikado211 12 months ago

That's the point here, technically all the options are good and will work, but since we are on a small amount of data Lambda will be the cheapest one, usually Glue or EMR will be kept for a big amount of data.

Here is a topic where people did a comparison in comments :

https://www.reddit.com/r/aws/comments/9umxv1/aws_glue_vs_lambda_costbenefit/

upvoted 4 times

TariqKipkemei 1 year, 1 month ago

Selected Answer: C

"processing will require 1 GB of memory and will finish within 30 seconds", perfect for AWS Lambda.

upvoted 4 times

antropaws 1 year, 6 months ago

Selected Answer: C

I reckon C, but I would consider other well founded options.

upvoted 3 times

nosense 1 year, 6 months ago

Selected Answer: C

c anyway the MOST cost-effectively

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #464

Topic 1

A company hosts an online shopping application that stores all orders in an Amazon RDS for PostgreSQL Single-AZ DB instance. Management wants to eliminate single points of failure and has asked a solutions architect to recommend an approach to minimize database downtime without requiring any changes to the application code.

Which solution meets these requirements?

- A. Convert the existing database instance to a Multi-AZ deployment by modifying the database instance and specifying the Multi-AZ option. **Most Voted**
- B. Create a new RDS Multi-AZ deployment. Take a snapshot of the current RDS instance and restore the new Multi-AZ deployment with the snapshot.
- C. Create a read-only replica of the PostgreSQL database in another Availability Zone. Use Amazon Route 53 weighted record sets to distribute requests across the databases.
- D. Place the RDS for PostgreSQL database in an Amazon EC2 Auto Scaling group with a minimum group size of two. Use Amazon Route 53 weighted record sets to distribute requests across instances.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: A

A. Convert the existing database instance to a Multi-AZ deployment by modifying the database instance and specifying the Multi-AZ option
upvoted 6 times

TariqKipkemei **Highly Voted** 1 year, 5 months ago

Selected Answer: A

Eliminate single points of failure = Multi-AZ deployment
upvoted 5 times

awsgeek75 **Most Recent** 10 months, 3 weeks ago

Selected Answer: A

A is correct but reason need to be clarified:

<https://aws.amazon.com/blogs/database/best-practices-for-converting-a-single-az-amazon-rds-instance-to-a-multi-az-instance/>

The instance doesn't automatically convert to Multi-AZ immediately. By default it will convert at next maintenance window but you can convert it immediately. Compared to B this is much better. CD are too many changes overall so unsuitable.

upvoted 3 times

antropaws 1 year, 6 months ago

Selected Answer: A

A)
<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZSingleStandby.html#Concepts.MultiAZ.Migrating>

upvoted 2 times

Abrar2022 1 year, 6 months ago

"minimize database downtime" so why create a new DB just modify the existing one so no time is wasted.

upvoted 5 times

cloudenthusiast 1 year, 6 months ago

Selected Answer: A

Compared to other solutions that involve creating new instances, restoring snapshots, or setting up replication manually, converting to a Multi-AZ deployment is a simpler and more streamlined approach with lower overhead.

Overall, option A offers a cost-effective and efficient way to minimize database downtime without requiring significant changes or additional complexities.

upvoted 3 times

Efren 1 year, 6 months ago

A for HA, but also read replica can convert itself to master if the master is down... so not sure if C?

upvoted 1 times

Efren 1 year, 6 months ago

Sorry, the Route 53 doesn't make sense to send requests to RR, what if it is a write?

upvoted 2 times

nosense 1 year, 6 months ago

Selected Answer: A

i guess aa

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #465

Topic 1

A company is developing an application to support customer demands. The company wants to deploy the application on multiple Amazon EC2 Nitro-based instances within the same Availability Zone. The company also wants to give the application the ability to write to multiple block storage volumes in multiple EC2 Nitro-based instances simultaneously to achieve higher application availability.

Which solution will meet these requirements?

- A. Use General Purpose SSD (gp3) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach
- B. Use Throughput Optimized HDD (st1) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach
- C. Use Provisioned IOPS SSD (io2) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach **Most Voted**
- D. Use General Purpose SSD (gp2) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach

Correct Answer: C*Community vote distribution***Comments****potomac** **Highly Voted** 1 year, 1 month ago**Selected Answer: C**

Multi-Attach is supported exclusively on Provisioned IOPS SSD (io1 and io2) volumes.

upvoted 12 times

awsgeek75 **Highly Voted** 10 months, 3 weeks ago**Selected Answer: C**

hdd<gp2<gp3<io2

upvoted 7 times

master9 **Most Recent** 11 months, 2 weeks ago**Selected Answer: C**

AWS IO2 does support Multi-Attach. Multi-Attach allows you to share access to an EBS data volume between up to 16 Nitro-based EC2 instances within the same Availability Zone. Each attached instance has full read and write permission to the shared volume. This feature is intended to make it easier to achieve higher application availability for customers that want to deploy

applications that manage storage consistency from multiple writers in shared storage infrastructure. However, please note that Multi-Attach on io2 is available in certain regions only.

upvoted 6 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: C

C. Use Provisioned IOPS SSD (io2) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach

upvoted 5 times

TariqKipkemei 1 year, 5 months ago

Selected Answer: C

Multi-Attach is supported exclusively on Provisioned IOPS SSD (io1 and io2) volumes.

[https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#:~:text=Multi%2DAttach%20is%20supported%20exclusively%20on%20Provisioned%20IOPS%20SSD%20\(io1%20and%20io2\)%20volumes.](https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#:~:text=Multi%2DAttach%20is%20supported%20exclusively%20on%20Provisioned%20IOPS%20SSD%20(io1%20and%20io2)%20volumes.)

upvoted 3 times

Axeashes 1 year, 5 months ago

Multi-Attach is supported exclusively on Provisioned IOPS SSD (io1 and io2) volumes.

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html>

upvoted 3 times

Uzi_m 1 year, 6 months ago

The correct answer is A.

Currently, Multi Attach EBS feature is supported by gp3 volumes also.

Multi-Attach is supported for certain EBS volume types, including io1, io2, gp3, st1, and sc1 volumes.

upvoted 2 times

Kp88 1 year, 4 months ago

No , Read this --> <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#considerations>

upvoted 4 times

AshishRocks 1 year, 6 months ago

Answer should be D

upvoted 1 times

Kp88 1 year, 4 months ago

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#considerations>

upvoted 2 times

AshishRocks 1 year, 6 months ago

By ChatGPT - Create General Purpose SSD (gp2) volumes: Provision multiple gp2 volumes with the required capacity for your application.

upvoted 1 times

Kp88 1 year, 4 months ago

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#considerations>

upvoted 1 times

AshishRocks 1 year, 6 months ago

Multi-Attach does not support Provisioned IOPS SSD (io2) volumes. Multi-Attach is currently available only for General Purpose SSD (gp2), Throughput Optimized HDD (st1), and Cold HDD (sc1) EBS volumes.

upvoted 1 times

Abrar2022 1 year, 6 months ago

Multi-Attach is supported exclusively on Provisioned IOPS SSD (io1 or io2) volumes.

upvoted 2 times

elmogy 1 year, 6 months ago

Selected Answer: C

only io1/io2 supports Multi-Attach
upvoted 3 times

Uzi_m 1 year, 6 months ago

Multi-Attach is supported for certain EBS volume types, including io1, io2, gp3, st1, and sc1 volumes.
upvoted 1 times

Kp88 1 year, 4 months ago

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html#considerations>
upvoted 2 times

examtopictempacc 1 year, 6 months ago

Selected Answer: C

only io1/io2 supports Multi-Attach
upvoted 3 times

Vlad 1 year, 6 months ago

Selected Answer: A

Option D suggests using General Purpose SSD (gp2) EBS volumes with Amazon EBS Multi-Attach. While gp2 volumes support multi-attach, gp3 volumes offer a more cost-effective solution with enhanced performance characteristics.
upvoted 1 times

Vlad 1 year, 6 months ago

I'm sorry :

Multi-Attach enabled volumes can be attached to up to 16 instances built on the Nitro System that are in the same Availability Zone. Multi-Attach is supported exclusively on Provisioned IOPS SSD (io1 or io2) volumes.
upvoted 3 times

Vlad 1 year, 6 months ago

The answer is C:
upvoted 2 times

EA100 1 year, 6 months ago

Answer - C
C. Use Provisioned IOPS SSD (io2) EBS volumes with Amazon Elastic Block Store (Amazon EBS) Multi-Attach.

While both option C and option D can support Amazon EBS Multi-Attach, using Provisioned IOPS SSD (io2) EBS volumes provides higher performance and lower latency compared to General Purpose SSD (gp2) volumes. This makes io2 volumes better suited for demanding and mission-critical applications where performance is crucial.

If the goal is to achieve higher application availability and ensure optimal performance, using Provisioned IOPS SSD (io2) EBS volumes with Multi-Attach will provide the best results.
upvoted 3 times

nosense 1 year, 6 months ago

Selected Answer: C

c is right
Amazon EBS Multi-Attach enables you to attach a single Provisioned IOPS SSD (io1 or io2) volume to multiple instances that are in the same Availability Zone.
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-volumes-multi.html>
nothing about gp
upvoted 3 times

cloudenthusiast 1 year, 6 months ago

Selected Answer: D

Given that the scenario does not mention any specific requirements for high-performance or specific IOPS needs, using General Purpose SSD (gp2) EBS volumes with Amazon EBS Multi-Attach (option D) is typically the more cost-effective and suitable choice. General Purpose SSD (gp2) volumes provide a good balance of performance and cost, making them well-suited for general-purpose workloads.

upvoted 1 times

elmogy 1 year, 6 months ago

the question has not mentioned anything about cost-effective solution.
only io1/io2 supports Multi-Attach

plus fyi, gp3 is the one gives a good balance of performance and cost. so gp2 is wrong in every way

upvoted 2 times

omoakin 1 year, 6 months ago

I agree

General Purpose SSD (gp2) volumes are the most common volume type. They were designed to be a cost-effective storage option for a wide variety of workloads. Gp2 volumes cover system volumes, dev and test environments, and various low-latency apps.

upvoted 1 times

y0 1 year, 6 months ago

gp2 - IOPS 16000

Nitro - IOPS 64000 - supported by io2. C is correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #466

Topic 1

A company designed a stateless two-tier application that uses Amazon EC2 in a single Availability Zone and an Amazon RDS Multi-AZ DB instance. New company management wants to ensure the application is highly available.

What should a solutions architect do to meet this requirement?

- A. Configure the application to use Multi-AZ EC2 Auto Scaling and create an Application Load Balancer **Most Voted**
- B. Configure the application to take snapshots of the EC2 instances and send them to a different AWS Region
- C. Configure the application to use Amazon Route 53 latency-based routing to feed requests to the application
- D. Configure Amazon Route 53 rules to handle incoming requests and create a Multi-AZ Application Load Balancer

Correct Answer: A

Community vote distribution

A (100%)

Comments

nosense **Highly Voted** 1 year ago

Selected Answer: A

it's A

upvoted 6 times

cloudenthusiast **Highly Voted** 1 year ago

Selected Answer: A

By combining Multi-AZ EC2 Auto Scaling and an Application Load Balancer, you achieve high availability for the EC2 instances hosting your stateless two-tier application.

upvoted 5 times

Guru4Cloud **Most Recent** 9 months, 2 weeks ago

Selected Answer: A

A. Configure the application to use Multi-AZ EC2 Auto Scaling and create an Application Load Balancer

upvoted 2 times

TariqKipkemei 11 months, 3 weeks ago

Selected Answer: A

Selected Answer: A

Highly available = Multi-AZ EC2 Auto Scaling and Application Load Balancer.
upvoted 3 times

antropaws 1 year ago

Selected Answer: A

Most likely A.
upvoted 2 times