

Guru4Cloud Most Recent 1 year, 3 months ago

Selected Answer: A

Add geographic restrictions to the content in CloudFront by using an allow list. Set up a custom error message
upvoted 3 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: A

Add geographic restrictions to the content in CloudFront by using an allow list. Set up a custom error message.
upvoted 2 times

jaydesai8 1 year, 4 months ago

Selected Answer: A

A makes sense - cloudfront has the capabilities of georestriction
upvoted 2 times

antropaws 1 year, 5 months ago

Selected Answer: A

Pretty sure it's A.
upvoted 2 times

alexandercamachop 1 year, 6 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/georestrictions.html>
upvoted 5 times

AncaZalog 1 year, 6 months ago

is B not A
upvoted 1 times

manuh 1 year, 5 months ago

Signed url or cookies can be used for the banner country as well?
upvoted 1 times

antropaws 1 year, 5 months ago

Why's that?
upvoted 2 times

awsgeek75 10 months, 3 weeks ago

How is signed URL going to be geo restricted? Anyone with signed url can access the content on that url regardless of their location so B is wrong.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #539

Topic 1

A company wants to use the AWS Cloud to improve its on-premises disaster recovery (DR) configuration. The company's core production business application uses Microsoft SQL Server Standard, which runs on a virtual machine (VM). The application has a recovery point objective (RPO) of 30 seconds or fewer and a recovery time objective (RTO) of 60 minutes. The DR solution needs to minimize costs wherever possible.

Which solution will meet these requirements?

- A. Configure a multi-site active/active setup between the on-premises server and AWS by using Microsoft SQL Server Enterprise with Always On availability groups.
- B. Configure a warm standby Amazon RDS for SQL Server database on AWS. Configure AWS Database Migration Service (AWS DMS) to use change data capture (CDC). **Most Voted**
- C. Use AWS Elastic Disaster Recovery configured to replicate disk changes to AWS as a pilot light.
- D. Use third-party backup software to capture backups every night. Store a secondary set of backups in Amazon S3.

Correct Answer: B*Community vote distribution*

B (59%)

C (41%)

Comments**1Alpha1** **Highly Voted** 1 year, 4 months ago**Selected Answer: B**

Backup & Restore (RPO in hours, RTO in 24 hours or less)
Pilot Light (RPO in minutes, RTO in hours)
Warm Standby (RPO in seconds, RTO in minutes) *** Right Answer ***
Active-Active (RPO is none or possibly seconds, RTO in seconds)
upvoted 10 times

1Alpha1 1 year, 4 months ago

[https://disaster-recovery.workshop.aws/en/intro/disaster-recovery.html#:~:text=Pilot%20Light%20\(RPO%20in%20minutes,that%20includes%20that%20critical%20core.](https://disaster-recovery.workshop.aws/en/intro/disaster-recovery.html#:~:text=Pilot%20Light%20(RPO%20in%20minutes,that%20includes%20that%20critical%20core.)
upvoted 4 times

pentium75 **Highly Voted** 1 year, 5 months ago

Selected Answer: C

Not A - too expensive and not using AWS services

Not B - "RDS for SQL Server" does not support everything that "SQL Server Standard which runs on a VM" does; CDC supports even less (https://docs.aws.amazon.com/dms/latest/userguide/CHAP_Source.SQLServer.html). Also it would be more expensive than C.

Not D - "Every night" would not meet the RPO requirement

upvoted 7 times

awsgeek75 1 year, 5 months ago

Thanks I was confused between B and C. This makes perfect sense!

upvoted 1 times

zdi561 **Most Recent** 4 months ago

Selected Answer: C

AWS Elastic Disaster Recovery (AWS DRS):

A service that can replicate entire virtual machines, including your SQL Server instance, to a secondary region for disaster recovery

upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

Selected Answer: B

AWS DMS with CDC can replicate changes from the on-premises SQL Server to the RDS instance in near real-time, meeting the 30-second RPO requirement.

A warm standby RDS instance ensures a faster recovery, meeting the 60-minute RTO requirement.

This solution is more cost-effective than upgrading to SQL Server Enterprise and provides a managed, scalable database service.

upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

C - Not optimized for real-time replication, making it challenging to meet the 30-second RPO requirement.

upvoted 1 times

Saliilgen 5 months, 2 weeks ago

Selected Answer: B

IMO answer is B.

AWS Elastic Disaster Recovery is cheaper and may meet the RTO/RPO requirements of the question, but option C states that it is configured as a pilot light (RTO/RPO 10s minute)

upvoted 2 times

LeonSauveterre 6 months ago

Selected Answer: C

RPO: the maximum acceptable amount of *data loss* during an outage or disaster.

RTO: the maximum acceptable amount of *downtime* after a failure.

Why B (RDS with DMS) might be better:

1. RTO Guarantee: A warm standby database provides faster recovery since the database is preconfigured and ready to accept traffic.
2. Simplicity: For workloads compatible with RDS and DMS, this approach simplifies management.

Why C (Elastic DR) might be better:

1. Workload Compatibility: If the on-premises SQL Server workload uses features that RDS or DMS cannot support, Elastic DR ensures full replication by mirroring the entire VM or disk.
2. Cost: Elastic DR replicates data but does not run a live database, making it much cheaper than a warm standby RDS.

"The DR solution needs to minimize costs wherever possible" - That makes me think C is the answer.

upvoted 2 times

FlyingHawk 4 months, 2 weeks ago

Elastic DRS is better suited for applications with less stringent RPO requirements. Achieving a 30-second RPO with this solution could be challenging, as it is not optimized for real-time replication.

upvoted 1 times

example_ 10 months, 2 weeks ago

Selected Answer: B

Pilot light (RPO in minutes, RTO in tens of minutes)

Warm standby (RPO in seconds, RTO in minutes)

https://docs.aws.amazon.com/wellarchitected/latest/reliability-pillar/rel_planning_for_recovery_disaster_recovery.html

upvoted 2 times

MrAliMohsan 9 months, 3 weeks ago

When cost is a concern, and you wish to achieve a similar RPO and RTO objectives as defined in the warm standby strategy, you could consider cloud native solutions, like AWS Elastic Disaster Recovery, that take the pilot light approach and offer improved RPO and RTO targets.

upvoted 3 times

abhiarns 1 year, 1 month ago

Selected Answer: C

AWS DRS(AWS Elastic Disaster Recovery) enables RPOs of seconds and RTOs of minutes.

upvoted 2 times

osmk 1 year, 3 months ago

Selected Answer: B

<https://docs.aws.amazon.com/whitepapers/latest/disaster-recovery-workloads-on-aws/disaster-recovery-options-in-the-cloud.html#warm-standby>

upvoted 3 times

awsgeek75 1 year, 4 months ago

Selected Answer: C

A: Not possible

B: With RDS it means your failover will launch a different database engine. This is wrong in general

D: No comments

C: It is a disk based replication so it will be similar DB server and this is the product managed by AWS for the DR of on-prem setups.

<https://aws.amazon.com/blogs/modernizing-with-aws/how-to-set-up-disaster-recovery-for-sql-server-always-on-availability-groups-using-aws-elastic-disaster-recovery/>

upvoted 2 times

1rob 1 year, 6 months ago

Selected Answer: C

AWS Elastic Disaster Recovery

If you are considering the pilot light or warm standby strategy for disaster recovery, AWS Elastic Disaster Recovery could provide an alternative approach with improved benefits. Elastic Disaster Recovery can offer an RPO and RTO target similar to warm standby, but maintain the low-cost approach of pilot light

From <https://docs.aws.amazon.com/wellarchitected/latest/reliability-pillar/rel_planning_for_recovery_disaster_recovery.html>

upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: B

With the pilot light approach, you replicate your data from one environment to another and provision a copy of your core workload infrastructure, not the fully functional copy of your production environment in a recovery environment.

upvoted 2 times

saymolet 1 year, 6 months ago

<https://aws.amazon.com/blogs/architecture/disaster-recovery-dr-architecture-on-aws-part-iii-pilot-light-and-warm-standby/>

upvoted 3 times

pentium75 1 year, 5 months ago

We have no idea if they are using SQL Server features that require OS customization etc., so we can't assume that the app would run on RDS for SQL Server at all. We need a replica of the VM that SQL Server is currently running on, thus C.
upvoted 2 times

thanhnv142 1 year, 7 months ago

C: Pilot light

- In pilot light, databases are always on, thus minimize RPO (can satisfy the 30s requirement)
 - Only apps are turn off. But it can satisfy the 60 minutes requirement
 - Warm standby, of course, can satisfy all the RPO and RTO requirements, but it is more expensive than pilot light
- upvoted 4 times

richguo 1 year, 8 months ago

Selected Answer: C

B(warm standby) is doable, but C (pilot light) is most cost effectively.
<https://aws.amazon.com/tw/blogs/architecture/disaster-recovery-dr-architecture-on-aws-part-iii-pilot-light-and-warm-standby/>
upvoted 3 times

LazyTs 1 year, 9 months ago

Selected Answer: B

The company wants to improve... so needs something guaranteed to be better than 60 mins RTO
upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

Configure a warm standby Amazon RDS for SQL Server database on AWS. Configure AWS Database Migration Service (AWS DMS) to use change data capture (CDC).
upvoted 3 times

Eminenza22 1 year, 9 months ago

Warm standby is costlier than Pilot Light
upvoted 3 times

PantryRaid 1 year, 9 months ago

Selected Answer: C

AWS DRS enables RPOs of seconds and RTOs of minutes. Pilot light is also cheaper than warm standby.
<https://aws.amazon.com/disaster-recovery/>
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #540

Topic 1


A company has an on-premises server that uses an Oracle database to process and store customer information. The company wants to use an AWS database service to achieve higher availability and to improve application performance. The company also wants to offload reporting from its primary database system.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. Use AWS Database Migration Service (AWS DMS) to create an Amazon RDS DB instance in multiple AWS Regions. Point the reporting functions toward a separate DB instance from the primary DB instance.
- B. Use Amazon RDS in a Single-AZ deployment to create an Oracle database. Create a read replica in the same zone as the primary DB instance. Direct the reporting functions to the read replica.
- C. Use Amazon RDS deployed in a Multi-AZ cluster deployment to create an Oracle database. Direct the reporting functions to use the reader instance in the cluster deployment.
- D. Use Amazon RDS deployed in a Multi-AZ instance deployment to create an Amazon Aurora database. Direct the reporting functions to the reader instances. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

mrsoa **Highly Voted** 1 year, 10 months ago

Selected Answer: D

Its D
Multi-AZ DB clusters aren't available with the following engines:
RDS for MariaDB
RDS for Oracle
RDS for SQL Server

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.RDS_Fea_Regions_DB-eng.Feature.MultiAZDBClusters.html
upvoted 35 times

alexandercamachop **Highly Voted** 2 years ago

Selected Answer: D

Selected Answer: C

C. Use Amazon RDS deployed in a Multi-AZ cluster deployment to create an Oracle database. Direct the reporting functions to use the reader instance in the cluster deployment.

A and B discarded.

The answer is between C and D

D says use an Amazon RDS to build an Amazon Aurora, makes no sense.

C is the correct one, high availability in multi az deployment.

Also point the reporting to the reader replica.

upvoted 13 times

1rob 1 year, 6 months ago

Multi-AZ DB clusters aren't available with the following engines:

RDS for MariaDB

RDS for Oracle

RDS for SQL Server

upvoted 6 times

bogobob 1 year, 6 months ago

using RDS to build Aurora from an Oracle DB <https://aws.amazon.com/tutorials/break-free-from-legacy-databases/migrate-oracle-to-amazon-aurora/>

upvoted 3 times

LeonSauveterre **Most Recent** 6 months ago

Selected Answer: D

In the MOST operationally efficient way, D is better, and A is too much. HOWEVER, Aurora is not Oracle, so this would require migrating from Oracle to Aurora (potentially requiring significant changes to the application, schema, and queries). The question did not say anything about not wanting to migrate.

B - 100% wrong. Single-AZ deployment does not meet the higher availability requirement.

C - 100% wrong. There's no way to create an Oracle db using RDS.

upvoted 1 times

chest_jd 7 months, 1 week ago

Unfortunately ALL ANSWERS ARE NOT CORRECT.

A. Use AWS Database Migration Service (AWS DMS) to create an Amazon RDS DB instance in multiple AWS Regions.
How can I create RDS using DMS ?

B. Single-Zone is not suitable for high available solution.

C. RDS cluster deployment are not available for DBs: RDS for Db2, RDS for MariaDB, RDS for Oracle, RDS for SQL Server
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.RDS_Fea_Regions_DB-eng.Feature.MultiAZDBClusters.html

D. Use Amazon RDS deployed in a Multi-AZ instance deployment to create an Amazon Aurora database.
Aurora DOES NOT support Oracle DB. Only MySQL and PostgreSQL

upvoted 1 times

Ravan 1 year, 3 months ago

Selected Answer: D

Multi-AZ (Availability Zone) deployments are not available for the following Amazon RDS database engines:

1. Amazon Aurora with MySQL compatibility
2. Amazon Aurora with PostgreSQL compatibility
3. Amazon RDS for SQL Server Express Edition
4. Amazon RDS for Oracle Standard Edition One
5. Amazon RDS for Oracle Standard Edition
6. Amazon RDS for Oracle SE2 (Standard Edition 2)

For these database engines, Amazon RDS provides high availability using other mechanisms specific to each engine, such as Read Replicas or different standby configurations. However, Multi-AZ deployments, which automatically provision and maintain a synchronous standby replica in a different Availability Zone for failover support, are not supported for these engines.

upvoted 4 times

noircesar25 1 year, 3 months ago

this link explains why the answer is C and confirms that RDS for Oracle supports multi-AZ
<https://aws.amazon.com/blogs/aws/multi-az-option-for-amazon-rds-oracle/>
upvoted 2 times

osmk 1 year, 3 months ago

Selected Answer: D

requiring high availability and performance. <https://aws.amazon.com/rds/aurora/>
upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: D

Between C&D, D is correct as C is not possible:
<https://aws.amazon.com/blogs/aws/multi-az-option-for-amazon-rds-oracle/>
upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: D

Not A - Creating multiple instances and keeping them in sync in DMS is surely not "operationally efficient"
Not B - "replica in the same zone" -> does not provide "higher availability"
Not C - "Multi-AZ cluster" does not support Oracle engine

Thus D. Question does not mention that the app would use Oracle-specific features; we're also not asked to minimize application changes. Ideal solution from AWS point of view is to move from Oracle to Aurora.
upvoted 2 times

XXXXXINN 8 months, 1 week ago

but Aurora doesn't support Oracle DB though
upvoted 1 time

aws94 1 year, 5 months ago

Selected Answer: C

i am sure just look here
<https://aws.amazon.com/ar/blogs/aws/amazon-rds-multi-az-db-cluster/>
upvoted 1 time

aws94 1 year, 5 months ago

sorry this is the right link:
<https://aws.amazon.com/ar/blogs/aws/multi-az-option-for-amazon-rds-oracle/>
upvoted 2 times

pentium75 1 year, 5 months ago

"Multi-AZ cluster (!)" does not support Oracle. Multi-AZ instance would.
upvoted 2 times

EEK2k 1 year, 6 months ago

Selected Answer: C

It should be C. Oracle DB is supported in RDS Multi-AZ with one standby for HA. <https://aws.amazon.com/rds/features/multi-az/>. Additionally, a reader instance/replica could be added to RDS Multi-AZ with one standby setup to offload the read requests. Aurora is only supported MySQL and Postgres compatible DB so "D" is out.
upvoted 2 times

1rob 1 year, 6 months ago

<https://aws.amazon.com/rds/features/multi-az/> gives: Amazon RDS Multi-AZ is available for RDS for PostgreSQL, RDS for MySQL, RDS for MariaDB, RDS for SQL Server, RDS for Oracle, and RDS for Db2. Amazon RDS Multi-AZ with two readable standbys is available for RDS for PostgreSQL and RDS for MySQL.
So no reader instance.

upvoted 1 times

potomac 1 year, 7 months ago

Selected Answer: D

Multi-AZ DB clusters are NOT available with the following engines:

RDS for MariaDB

RDS for Oracle

RDS for SQL Server

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.RDS_Fea_Regions_DB-eng.Feature.MultiAZDBClusters.html

upvoted 2 times

danielmakita 1 year, 7 months ago

It is C. Aurora database doesn't support Oracle.

upvoted 1 times

JA2018 6 months, 2 weeks ago

as pentium75 had pointed out, the stem does not mention that the app would use Oracle-specific features; we're also not asked to minimize application changes. Ideal solution from AWS point of view is to move from Oracle to Aurora.

upvoted 1 times

wsdasdasdqwda 1 year, 7 months ago

You can use Aurora instead of Oracle. There are tutorials how to migrate Oracle to Aurora. On top C is not supported. The is not Multi-AZ DB CLUSTER for Oracle.

upvoted 3 times

wsdasdasdqwda 1 year, 7 months ago

It is D

upvoted 1 times

thanhnv142 1 year, 7 months ago

None options seems valid. Not C because it is not supported. But not D as well. RDS is not Aurora. They are two separate services. Additionally, In multi AZ instance deployment, it only provides fault tolerance, not High avail.

upvoted 2 times

Nikki013 1 year, 9 months ago

Selected Answer: D

Multi-AZ Cluster does not support Oracle as engine:

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.RDS_Fea_Regions_DB-eng.Feature.MultiAZDBClusters.html

upvoted 2 times

Bennyboy789 1 year, 9 months ago

Selected Answer: D

D is my choice.

Multi-AZ DB cluster does not support Oracle DB.

upvoted 3 times

rjbihari 1 year, 9 months ago

Option C is correct one .

As there is no option for 'Aurora(Oracle Compatible)'.so this kick out D from race.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #541

Topic 1

A company wants to build a web application on AWS. Client access requests to the website are not predictable and can be idle for a long time. Only customers who have paid a subscription fee can have the ability to sign in and use the web application.

Which combination of steps will meet these requirements MOST cost-effectively? (Choose three.)

A. Create an AWS Lambda function to retrieve user information from Amazon DynamoDB. Create an Amazon API Gateway endpoint to accept RESTful APIs. Send the API calls to the Lambda function. **Most Voted**

B. Create an Amazon Elastic Container Service (Amazon ECS) service behind an Application Load Balancer to retrieve user information from Amazon RDS. Create an Amazon API Gateway endpoint to accept RESTful APIs. Send the API calls to the Lambda function.

C. Create an Amazon Cognito user pool to authenticate users. **Most Voted**


D. Create an Amazon Cognito identity pool to authenticate users.

E. Use AWS Amplify to serve the frontend web content with HTML, CSS, and JS. Use an integrated Amazon CloudFront configuration. **Most Voted**

F. Use Amazon S3 static web hosting with PHP, CSS, and JS. Use Amazon CloudFront to serve the frontend web content.

Correct Answer: ACE

Community vote distribution

**Comments**

manOfThePeople **Highly Voted** 1 year, 9 months ago

If in doubt between E or F. S3 doesn't support server-side scripts, PHP is a server-side script. The answer is ACE.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteHosting.html>

upvoted 21 times

msdnpro **Highly Voted** 1 year, 11 months ago

Selected Answer: ACE

Option B (Amazon ECS) is not the best option since the website "can be idle for a long time", so Lambda (Option A) is a more cost-effective choice.

Option D is incorrect because User pools are for authentication (identity verification) while Identity pools are for authorization (access control).

Option F is wrong because S3 web hosting only supports static web files like HTML/CSS, and does not support PHP or JavaScript.

upvoted 10 times

0628atv 1 year, 11 months ago

https://aws.amazon.com/getting-started/projects/build-serverless-web-app-lambda-apigateway-s3-dynamodb-cognito/module-1/?nc1=h_ls

upvoted 3 times

Gape4 Most Recent 11 months, 2 weeks ago

Selected Answer: ACE

I will go for A C E

upvoted 2 times

awsgeek75 1 year, 5 months ago

Selected Answer: ACE

A: App may be idle for long time so Lambda is perfect (charge per invocation)

C: Cognito user pool for user auth

E: Amplify is low code web dev tool

B: Wrong, too much cost when idle

D: Identity pool is session management/identification. Does not help with auth.

F: S3 + PHP doesn't work also no security

upvoted 3 times

rcptryk 1 year, 6 months ago

Selected Answer: ACE

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteHosting.html>

S3 doesn't support server-side scripting

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: ACE

User Pool = authentication

Identity Pool = authorization

upvoted 7 times

thanhnv142 1 year, 7 months ago

A D F:

A: for hosting the dynamic content of the app. Pay as execution

D: for granting temporary privilege access to users who has paid a fee.

F: for hosting the static content of the app

upvoted 1 times

awsgeek75 1 year, 4 months ago

There is no static content in this web application so F is wrong. You cannot host PHP on S3 also so it is just wrong.

upvoted 2 times

kwang312 1 year, 8 months ago

Selected Answer: ACE

ACE is correct answer

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: CEF

C) Create an Amazon Cognito user pool to authenticate users

C) Create an Amazon Cognito user pool to authenticate users.

E) Use AWS Amplify to serve the frontend web content with HTML, CSS, and JS. Use an integrated CloudFront configuration.

F) Use Amazon S3 static web hosting with PHP, CSS, and JS. Use Amazon CloudFront to serve the frontend web content.

upvoted 1 times

awsgeek75 1 year, 4 months ago

There is no static content in this web application so F is wrong. You cannot host PHP on S3 also so it is just wrong.

upvoted 2 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: ACE

Build a web application = AWS Amplify

Sign in users = Amazon Cognito user pool

Traffic can be idle for a long time = AWS Lambda

Amazon S3 does not support server-side scripting such as PHP, JSP, or ASP.NET.

[https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteHosting.html?](https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteHosting.html?icmpid=docs_amazons3_console#:~:text=website%20relies%20on,-server%2Dside,-processing%2C%20including%20server)

[icmpid=docs_amazons3_console#:~:text=website%20relies%20on,-server%2Dside,-processing%2C%20including%20server](https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteHosting.html?icmpid=docs_amazons3_console#:~:text=website%20relies%20on,-server%2Dside,-processing%2C%20including%20server)

Traffic can be idle for a long time = AWS Lambda

upvoted 2 times

james2033 1 year, 10 months ago

Selected Answer: ACE

Use exclusion method: No need for Container (no need run all time), remove B. PHP cannot run with static Amazon S3, remove F.

Use selection method: Idle for sometime, choose AWS Lambda, choose A. "Amazon Cognito is an identity platform for web and mobile apps." (<https://docs.aws.amazon.com/cognito/latest/developerguide/what-is-amazon-cognito.html>), choose C.

Create an identity pool <https://docs.aws.amazon.com/cognito/latest/developerguide/tutorial-create-identity-pool.html> . AWS Amplify <https://aws.amazon.com/amplify/> for build full-stack web-app in hours.

upvoted 6 times

baba365 1 year, 11 months ago

Ans: ACE

use AWS SDK for PHP/JS with S3

https://docs.aws.amazon.com/sdk-for-php/v3/developer-guide/php_s3_code_examples.html

upvoted 1 times

unbendable 1 year, 7 months ago

did you actually read the link or just copy the first link from google here? the sdk is intended for usage in a php application. it does not say anything about php support in a s3 bucket

upvoted 2 times

Zox42 1 year, 11 months ago

Selected Answer: ACE

Answer is ACE

upvoted 2 times

jaydesai8 1 year, 11 months ago

Selected Answer: ACE

Lambda =serverless

User Pool = For user authentication

Amplify = hosting web/mobile apps

upvoted 3 times

live_reply_developers 1 year, 11 months ago

Selected Answer: ACE

S3 doesn't support PHP as stated in answer F.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/WebsiteHosting.html>
upvoted 4 times

wRhIH 1 year, 11 months ago

Selected Answer: ACE

I don't think S3 can handle anything dynamic such as PHP. So I go for ACE
upvoted 2 times

antropaws 1 year, 11 months ago

Selected Answer: ACF

ACF no doubt. Check the difference between user pools and identity pools.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #542

Topic 1

A media company uses an Amazon CloudFront distribution to deliver content over the internet. The company wants only premium customers to have access to the media streams and file content. The company stores all content in an Amazon S3 bucket. The company also delivers content on demand to customers for a specific purpose, such as movie rentals or music downloads.

Which solution will meet these requirements?

- A. Generate and provide S3 signed cookies to premium customers.
- B. Generate and provide CloudFront signed URLs to premium customers. **Most Voted**
- C. Use origin access control (OAC) to limit the access of non-premium customers.
- D. Generate and activate field-level encryption to block non-premium customers.

Correct Answer: B

Community vote distribution

B (100%)

Comments

NayeraB **Highly Voted** 9 months, 3 weeks ago

This question page is filled with premium customers I just can't upvoted 8 times

awsgeek75 **Most Recent** 11 months ago

Selected Answer: B

CloudFront Signed URL with Custom Policy are exactly for this.

A: Nope, cookies don't help as they don't restrict URL

C: Wrong. OAC for non-premium customers, how is that even possible without any details here?

D: Field encryption, while good idea, does not help restricting the content by customer

upvoted 3 times

pentium75 11 months, 1 week ago

Selected Answer: B

Authentication is done by Cloudfront, thus B

upvoted 4 times

ferdzcruz 11 months, 2 weeks ago

Content on demand = CloudFront. B

upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Generate and provide CloudFront signed URLs to premium customers.

upvoted 3 times

TariqKipkemei 1 year, 4 months ago

Selected Answer: B

Use CloudFront signed URLs or signed cookies to restrict access to documents, business data, media streams, or content that is intended for selected users, for example, users who have paid a fee.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/PrivateContent.html#:~:text=CloudFront%20signed%20URLs>

upvoted 3 times

james2033 1 year, 4 months ago

Selected Answer: B

See <https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/private-content-signed-urls.html#private-content-how-signed-urls-work>

upvoted 2 times

haoAWS 1 year, 5 months ago

Selected Answer: B

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/PrivateContent.html>

Notice that A is not correct because it should be CloudFront signed URL, not S3.

upvoted 3 times

antropaws 1 year, 5 months ago

Why not C?

upvoted 1 times

pentium75 11 months, 1 week ago

OAC requires the consumers to have an IAM role with access to the S3 content, this is not what we're after here.

upvoted 2 times

antropaws 1 year, 5 months ago

<https://aws.amazon.com/blogs/networking-and-content-delivery/amazon-cloudfront-introduces-origin-access-control-oac/>

upvoted 2 times

alexandercamachop 1 year, 6 months ago

Selected Answer: B

Signed URLs

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/PrivateContent.html>

upvoted 3 times

haoAWS 1 year, 5 months ago

Then why A is incorrect?

upvoted 1 times

pentium75 11 months, 1 week ago

Because Authentication is done by Cloudfront, not S3.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #543

Topic 1

A company runs Amazon EC2 instances in multiple AWS accounts that are individually billed. The company recently purchased a Savings Plan. Because of changes in the company's business requirements, the company has decommissioned a large number of EC2 instances. The company wants to use its Savings Plan discounts on its other AWS accounts.

Which combination of steps will meet these requirements? (Choose two.)

A. From the AWS Account Management Console of the management account, turn on discount sharing from the billing preferences section. **Most Voted**

B. From the AWS Account Management Console of the account that purchased the existing Savings Plan, turn on discount sharing from the billing preferences section. Include all accounts.


C. From the AWS Organizations management account, use AWS Resource Access Manager (AWS RAM) to share the Savings Plan with other accounts.

D. Create an organization in AWS Organizations in a new payer account. Invite the other AWS accounts to join the organization from the management account. **Most Voted**

E. Create an organization in AWS Organizations in the existing AWS account with the existing EC2 instances and Savings Plan. Invite the other AWS accounts to join the organization from the management account.

Correct Answer: AD

Community vote distribution



Comments

Agerim2010 **Highly Voted** 1 year, 11 months ago

i had this question today
upvoted 20 times

Nikki013 **Highly Voted** 1 year, 9 months ago

Selected Answer: AD

It is not recommended to have workload on the management account.
upvoted 6 times

JA2018 6 months, 2 weeks ago

confirm will fail the AWS WAR process
upvoted 1 times

zdi561 **Most Recent** 4 months, 2 weeks ago

Selected Answer: AD

google 'aws how to share ec2 shaving plans among accounts of an organization'
upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

Selected Answer: AE

A E, firstly we create an organizatio in the existing AWS account with savings plan, so this account becomes a management account. Secondly, from the console of this account, turn on discount sharing.
upvoted 1 times

MaxMingxing 5 months, 1 week ago

Selected Answer: AE

Answer from ChatGPT
upvoted 1 times

FlyingHawk 5 months, 3 weeks ago

Selected Answer: AB

If the organization is already created, I am feeling A and B are the right one as "To share an Reserved Instance or Savings Plans discount with an account, both accounts must have sharing activated."
<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html>. If no AWS organization was there, then you need to create an organization, then invite others join, D. then A and B.
upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

B is not correct, Discount sharing can only be enabled from the management account of an AWS Organization, not from the account where the Savings Plan was purchased.
The correct answer is A E.
<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html#ri-turn-on-process>
upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: AD

For me, E makes no sense as the discount is with a new payer and cannot be transferred to an existing account unless customer service is involved.
upvoted 3 times

awsgeek75 1 year, 5 months ago

Also, "A company runs Amazon EC2 instances in multiple AWS accounts that are individually bled"

It's not bled, it is "billed"
upvoted 6 times

JA2018 6 months, 2 weeks ago

ROTFL :P
upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: AD

Organization should be created by a new account that is reserved for management. Thus D, followed by A (discount sharing must be enabled in the management account).
upvoted 5 times

ErnShm 1 year, 9 months ago

LinSinn 1 year, 2 months ago

AE

<https://repost.aws/questions/QUQoJuQLNOTDiyEuCLARIBFQ/transfer-savings-plan-across-organizations#:~:text=AWS%20Support%20can%20transfer%20Savings%20Plans%20from%20the%20management%20account%20to%20a%20member%20account%20or%20from%20a%20member%20account%20to%20the%20management%20account%20within%20a%20single%20Organization%20with%20an%20AWS%20Support%20Case.>

upvoted 1 times

lemur88 1 year, 9 months ago

Selected Answer: AD

Not E - it mentions using an account with existing EC2s as the management account, which goes against the best practice for a management account

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_best-practices_mgmt-acct.html

upvoted 4 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: AE

AE is best

upvoted 1 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: AE

AE is best

upvoted 1 times

james2033 1 year, 10 months ago

Selected Answer: AE

- B is not accepted, because "include all accounts", remove B.
- D has "Create an organization in AWS Organization in a new payer account", it is wrong, remove D.
- at C: AWS Resource Access Manager (AWS RAM) <https://aws.amazon.com/ram/> it is for security, not for billing. Remove C.
Has A, E remain, and choosed.

A. "turn on discount sharing" is ok. This case: Has discount for many EC2 instances in one account, then want to share with other user. At E, create Organization, then share.

upvoted 2 times

pentium75 1 year, 5 months ago

What is the problem with "include all accounts"?

upvoted 2 times

antropaws 1 year, 11 months ago

Selected Answer: AE

I vote AE.

upvoted 2 times

MrAWSAssociate 1 year, 11 months ago

Selected Answer: AE

AE are correct !

upvoted 2 times

oras2023 1 year, 11 months ago

Selected Answer: CD

It's not good practice to create a payer account with any workload so it must be D.
By the reason that we need Organizations for sharing, then we need to turn on its from our PAYER account. (all sub-accounts start share discounts)

upvoted 1 times

oras2023 1 year, 11 months ago

changed to AD
upvoted 4 times

maver144 1 year, 12 months ago

Selected Answer: AE

@alexandercamachop it is AE. I believe its just typo. RAM is not needed anyhow.
upvoted 5 times

oras2023 1 year, 11 months ago

You are right
<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/ri-turn-off.html>
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #544

Topic 1

A retail company uses a regional Amazon API Gateway API for its public REST APIs. The API Gateway endpoint is a custom domain name that points to an Amazon Route 53 alias record. A solutions architect needs to create a solution that has minimal effects on customers and minimal data loss to release the new version of APIs.

Which solution will meet these requirements?

A. Create a canary release deployment stage for API Gateway. Deploy the latest API version. Point an appropriate percentage of traffic to the canary stage. After API verification, promote the canary stage to the production stage.

Most Voted

B. Create a new API Gateway endpoint with a new version of the API in OpenAPI YAML file format. Use the import-to-update operation in merge mode into the API in API Gateway. Deploy the new version of the API to the production stage.

C. Create a new API Gateway endpoint with a new version of the API in OpenAPI JSON file format. Use the import-to-update operation in overwrite mode into the API in API Gateway. Deploy the new version of the API to the production stage.

D. Create a new API Gateway endpoint with new versions of the API definitions. Create a custom domain name for the new API Gateway API. Point the Route 53 alias record to the new API Gateway API custom domain name.

Correct Answer: A

Community vote distribution

A (100%)

Comments

AudreyNguyenHN Highly Voted 1 year, 10 months ago

We made it all the way here. Good luck everyone!
upvoted 18 times

ddddddddddww12 Highly Voted 1 year, 11 months ago

what are the total number of questions this package has as on 14 July 2023 , is it 544 or 551 ?
upvoted 8 times

MrAliMohsan 9 months, 3 weeks ago

August 2024 - 981 questions :(

upvoted 3 times

Faridtnx 1 year, 2 months ago

March 2024 its 825 questions. Constantly adding.

Doe ur question, ExamTopic always shows a few more question in listing compared to actual number

upvoted 5 times

NayeraB 1 year, 3 months ago

It's 20th of Feb 2024, and it's 798 (it says 804 at the top I donno why tho)

upvoted 3 times

rosanna Most Recent 6 months ago

Selected Answer: A

where do you guys get to know the answer to these things? I study from Stephan's content but lots of topics are uncovered, where am I supposed to find all the exam content in one place?

upvoted 1 times

awsgeek75 1 year, 5 months ago

Selected Answer: A

<https://docs.aws.amazon.com/apigateway/latest/developerguide/canary-release.html>

upvoted 4 times

potomac 1 year, 7 months ago

Selected Answer: A

In a canary release deployment, total API traffic is separated at random into a production release and a canary release with a pre-configured ratio. Typically, the canary release receives a small percentage of API traffic and the production release takes up the rest. The updated API features are only visible to API traffic through the canary. You can adjust the canary traffic percentage to optimize test coverage or performance.

<https://docs.aws.amazon.com/apigateway/latest/developerguide/canary-release.html>

upvoted 7 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: A

Using a canary release deployment allows incremental rollout of the new API version to a percentage of traffic. This minimizes impact on customers and potential data loss during the release.

upvoted 3 times

TariqKipkemei 1 year, 10 months ago

Selected Answer: A

Minimal effects on customers and minimal data loss = Canary deployment

upvoted 5 times

james2033 1 year, 10 months ago

Selected Answer: A

Key word "canary release". See this term in See: <https://www.jetbrains.com/teamcity/ci-cd-guide/concepts/canary-release/> and/or <https://martinfowler.com/bliki/CanaryRelease.html>

upvoted 2 times

Abrar2022 1 year, 11 months ago

Selected Answer: A

keyword: "latest versions on an api"

Canary release is a software development strategy in which a "new version of an API" (as well as other software) is deployed for testing purposes.

upvoted 4 times

jkhan2405 1 year, 12 months ago

Selected Answer: A

It's A

upvoted 2 times

alexandercamachop 2 years ago

Selected Answer: A

A. Create a canary release deployment stage for API Gateway. Deploy the latest API version. Point an appropriate percentage of traffic to the canary stage. After API verification, promote the canary stage to the production stage.

Canary release meaning only certain percentage of the users.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #545

Topic 1


A company wants to direct its users to a backup static error page if the company's primary website is unavailable. The primary website's DNS records are hosted in Amazon Route 53. The domain is pointing to an Application Load Balancer (ALB). The company needs a solution that minimizes changes and infrastructure overhead.

Which solution will meet these requirements?

- A. Update the Route 53 records to use a latency routing policy. Add a static error page that is hosted in an Amazon S3 bucket to the records so that the traffic is sent to the most responsive endpoints.
- B. Set up a Route 53 active-passive failover configuration. Direct traffic to a static error page that is hosted in an Amazon S3 bucket when Route 53 health checks determine that the ALB endpoint is unhealthy. **Most Voted**
- C. Set up a Route 53 active-active configuration with the ALB and an Amazon EC2 instance that hosts a static error page as endpoints. Configure Route 53 to send requests to the instance only if the health checks fail for the ALB.
- D. Update the Route 53 records to use a multivalue answer routing policy. Create a health check. Direct traffic to the website if the health check passes. Direct traffic to a static error page that is hosted in Amazon S3 if the health check does not pass.

Correct Answer: B

Community vote distribution



Comments

TariqKipkemei **Highly Voted** 1 year, 6 months ago

Selected Answer: B

Set up a Route 53 active-passive failover configuration. Direct traffic to a static error page that is hosted in an Amazon S3 bucket when Route 53 health checks determine that the ALB endpoint is unhealthy.

upvoted 6 times

ssa03 **Highly Voted** 1 year, 9 months ago

Selected Answer: B

B is correct

upvoted 5 times

Danilus Most Recent 7 months ago

key -minimizes changes and infrastructure overhead

key- backup static error page

Option B is the appropriate solution because it sets up automatic failover to an alternative endpoint (the static error page in S3) when the ALB becomes unavailable. Option D is not viable as it does not implement a true failover mechanism and may continue to direct traffic to an unhealthy ALB.

upvoted 1 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

Setting up a Route 53 active-passive failover configuration with the ALB as the primary endpoint and an Amazon S3 static website as the passive endpoint meets the requirements with minimal overhead.

Route 53 health checks can monitor the ALB health. If the ALB becomes unhealthy, traffic will automatically failover to the S3 static website. This provides automatic failover with minimal configuration changes

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Sorry. I mean B

upvoted 6 times

Nirav1112 1 year, 10 months ago

B is correct

upvoted 3 times

mrsoa 1 year, 10 months ago

Selected Answer: B

B seems correct

upvoted 4 times

Bmaster 1 year, 10 months ago

B is correct..

<https://repost.aws/knowledge-center/fail-over-s3-r53>

upvoted 4 times

awsgeek75 1 year, 5 months ago

Nice link find!

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #546

Topic 1

A recent analysis of a company's IT expenses highlights the need to reduce backup costs. The company's chief information officer wants to simplify the on-premises backup infrastructure and reduce costs by eliminating the use of physical backup tapes. The company must preserve the existing investment in the on-premises backup applications and workflows.

What should a solutions architect recommend?

- A. Set up AWS Storage Gateway to connect with the backup applications using the NFS interface.
- B. Set up an Amazon EFS file system that connects with the backup applications using the NFS interface.
- C. Set up an Amazon EFS file system that connects with the backup applications using the iSCSI interface.
- D. Set up AWS Storage Gateway to connect with the backup applications using the iSCSI-virtual tape library (VTL) interface.

Most Voted**Correct Answer: D***Community vote distribution*

D (100%)

Comments**awsgeek75** **Highly Voted** 11 months ago**Selected Answer: D**

Tape... lol

The company must preserve its existing investment so they want to keep using existing applications. This means EFS won't work, and NFS may not be compatible. VTL is the only thing that may be compatible with an application workflow that backs up to tapes.

Who the hell comes up with these questions!

upvoted 6 times

TariqKipkemei **Highly Voted** 1 year ago**Selected Answer: D**

Use Tape Gateway to replace physical tapes on premises with virtual tapes on AWS—reducing your data storage costs without changing your tape-based backup workflows. Tape Gateway supports all leading backup applications and caches virtual tapes on premises for low-latency data access. It compresses your tape data, encrypts it, and stores it in a virtual tape library in Amazon Simple Storage Service (Amazon S3). From there, you can transfer it to either Amazon S3 Glacier Flexible Retrieval or

Amazon Simple Storage Service (Amazon S3). From there, you can transfer it to either Amazon S3 Glacier Flexible Retrieval or Amazon S3 Glacier Deep Archive to help minimize your long-term storage costs.

<https://aws.amazon.com/storagegateway/vtl/#:~:text=Use-,Tape%20Gateway,-to%20replace%20physical>
upvoted 5 times

Nisarg2121 Most Recent 1 year, 1 month ago

Selected Answer: D

Tape Gateway is use for attache with app.
upvoted 4 times

gouranga45 1 year, 1 month ago

Selected Answer: D

Option says it all
upvoted 4 times

Po_chih 1 year, 1 month ago

Selected Answer: D

Tape Gateway enables you to replace using physical tapes on premises with virtual tapes in AWS without changing existing backup workflows. Tape Gateway supports all leading backup applications and caches virtual tapes on premises for low-latency data access. Tape Gateway encrypts data between the gateway and AWS for secure data transfer, and compresses data and transitions virtual tapes between Amazon S3 and Amazon S3 Glacier Flexible Retrieval, or Amazon S3 Glacier Deep Archive, to minimize storage costs.

upvoted 3 times

ssa03 1 year, 3 months ago

Selected Answer: D

https://aws.amazon.com/storagegateway/vtl/?nc1=h_ls
upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

Set up AWS Storage Gateway to connect with the backup applications using the iSCSI-virtual tape library (VTL) interface.
upvoted 2 times

Bmaster 1 year, 4 months ago

D is correct

https://aws.amazon.com/storagegateway/vtl/?nc1=h_ls
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #547

Topic 1

A company has data collection sensors at different locations. The data collection sensors stream a high volume of data to the company. The company wants to design a platform on AWS to ingest and process high-volume streaming data. The solution must be scalable and support data collection in near real time. The company must store the data in Amazon S3 for future reporting.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon Kinesis Data Firehose to deliver streaming data to Amazon S3. **Most Voted**
- B. Use AWS Glue to deliver streaming data to Amazon S3.
- C. Use AWS Lambda to deliver streaming data and store the data to Amazon S3.
- D. Use AWS Database Migration Service (AWS DMS) to deliver streaming data to Amazon S3.

Correct Answer: A*Community vote distribution***Comments****emakid** 11 months, 2 weeks ago**Selected Answer: A**

Amazon Kinesis Data Firehose is a fully managed service for delivering real-time streaming data to destinations such as Amazon S3, Amazon Redshift, Amazon Elasticsearch Service, and Splunk. It requires minimal setup and maintenance, automatically scales to match the throughput of your data, and offers near real-time data delivery with minimal operational overhead.

upvoted 3 times

awsgeek75 1 year, 5 months ago**Selected Answer: A**

High volume streaming data = Kinesis
B: Glue is for ETL (to S3 is ok) but not for streaming
C: Lambda more overhead
D: Streaming != Data migration
upvoted 3 times

nentium75 1 year, 5 months ago

pentium75 1 year, 6 months ago

Selected Answer: A

sensor data = Kinesis
upvoted 4 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

Amazon Kinesis Data Firehose: Capture, transform, and load data streams into AWS data stores (S3) in near real-time.

https://aws.amazon.com/pm/kinesis/?gclid=CjwKCAiAu9yqBhBmEiwAHTx5px9z182o0HBEX0BGXU7VeOCOdNpkJMxgbSfvcHINKN4NHVnbEa0Y1xoCuU0QAvD_BwE&trk=239a97c0-9c5d-42a5-ac65-7381b62f3756&sc_channel=ps&ef_id=CjwKCAiAu9yqBhBmEiwAHTx5px9z182o0HBEX0BGXU7VeOCOdNpkJMxgbSfvcHINKN4NHVnbEa0Y1xoCuU0QAvD_BwE:G:s&s_kwid=AL!4422!3!651612444428!e!!g!!kinesis%20firehose!19836376048!149982297311#:~:text=Kinesis%20Data%20Firehose-,Capture%2C,-transform%2C%20and%20load
upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: A

A for sure
upvoted 3 times

ssa03 1 year, 9 months ago

Selected Answer: A

Correct Answer: A
upvoted 3 times

manOfThePeople 1 year, 9 months ago

A is the answer, near real-time = Kinesis Data Firehose.
upvoted 4 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: D

Use Amazon Kinesis Data Firehose to deliver streaming data to Amazon S3
upvoted 3 times

pentium75 1 year, 5 months ago

That is A
upvoted 2 times

bjexamprep 1 year, 9 months ago

Selected Answer: D

Kinesis Data Firehose is only real-time answer
upvoted 3 times

pentium75 1 year, 5 months ago

That is A
upvoted 3 times

mrsoa 1 year, 10 months ago

Selected Answer: A

A is the correct answer
upvoted 3 times

Deepakin96 1 year, 10 months ago

Selected Answer: A

Kinesis = Near Real Time

KINESIS – NEAR REAL TIME

upvoted 4 times

Kaiden123 1 year, 10 months ago

Selected Answer: A

Data collection in near real time = Amazon Kinesis Data Firehose

upvoted 4 times

Bmaster 1 year, 10 months ago

A is correct..

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #548

Topic 1


A company has separate AWS accounts for its finance, data analytics, and development departments. Because of costs and security concerns, the company wants to control which services each AWS account can use.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use AWS Systems Manager templates to control which AWS services each department can use.
- B. Create organization units (OUs) for each department in AWS Organizations. Attach service control policies (SCPs) to the OUs. **Most Voted**
- C. Use AWS CloudFormation to automatically provision only the AWS services that each department can use.
- D. Set up a list of products in AWS Service Catalog in the AWS accounts to manage and control the usage of specific AWS services.

Correct Answer: B

Community vote distribution



Comments

zdi561 4 months ago

Selected Answer: D

SCP set a guardrail (what you can not do, permission on IAM user and role), ASC provides templates services that you can only use.

upvoted 1 times

awsgeek75 11 months ago

Selected Answer: B

Departments = Organizational Units

upvoted 3 times

TariqKipkemei 1 year ago

Selected Answer: B

Create organization units (OUs) for each department in AWS Organizations. Attach service control policies (SCPs) to the OUs

upvoted 2 times

ssa03 1 year, 3 months ago

Selected Answer: B

Correct Answer: B

upvoted 2 times

lemur88 1 year, 3 months ago

Selected Answer: B

SCPs to centralize permissioning

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Create organization units (OUs) for each department in AWS Organizations. Attach service control policies (SCPs) to the OUs.

upvoted 2 times

xyb 1 year, 4 months ago

Selected Answer: B

control services --> SCP

upvoted 2 times

Ale1973 1 year, 4 months ago

Selected Answer: D

My rational: Scenary is "A company has separate AWS accounts", it is not mentioning anything about use of Organizations or needs related to centralized management of these accounts.

Then, set up a list of products in AWS Service Catalog in the AWS accounts (on each AWS account) is the best way to manage and control the usage of specific AWS services.

upvoted 1 times

pentium75 11 months, 1 week ago

"Separate AWS accounts" just says that it's multiple accounts, it does not indicate that they are NOT connected into a organization.

Service Catalog alone does not restrict anything. You'd need to create a service in Service Catalog for everything you're allowing to use, then grant permissions on those services, and you'd need to remove other permissions from everyone. All of which is not mentioned in D. Just "setting up a list of products in AWS Service Catalog in the AWS accounts" will not restrict anyone from doing what he could do before.

upvoted 3 times

mrsoa 1 year, 4 months ago

Selected Answer: B

BBBBBBBBBB

upvoted 2 times

Deepakin96 1 year, 4 months ago

Selected Answer: B

To control different AWS account you required AWS Organisation

upvoted 2 times

Bmaster 1 year, 4 months ago

B is correct!!!!

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #549

Topic 1

A company has created a multi-tier application for its ecommerce website. The website uses an Application Load Balancer that resides in the public subnets, a web tier in the public subnets, and a MySQL cluster hosted on Amazon EC2 instances in the private subnets. The MySQL database needs to retrieve product catalog and pricing information that is hosted on the internet by a third-party provider. A solutions architect must devise a strategy that maximizes security without increasing operational overhead.

What should the solutions architect do to meet these requirements?

- A. Deploy a NAT instance in the VPC. Route all the internet-based traffic through the NAT instance.
- B. Deploy a NAT gateway in the public subnets. Modify the private subnet route table to direct all internet-bound traffic to the NAT gateway. **Most Voted**
- C. Configure an internet gateway and attach it to the VPC. Modify the private subnet route table to direct internet-bound traffic to the internet gateway.
- D. Configure a virtual private gateway and attach it to the VPC. Modify the private subnet route table to direct internet-bound traffic to the virtual private gateway.

Correct Answer: B

Community vote distribution

B (100%)

Comments

awsgeek75 11 months ago

Selected Answer: B

A: Probably an old question so this option is here but NAT instance is overhead
 C: Not secure as IG opens up a lot of things
 D: VPG connects to a service
 B: NG is managed solution. Secure by config
 upvoted 3 times

TariqKipkemei 1 year ago

Selected Answer: B

Deploy a NAT gateway in the public subnets. Modify the private subnet route table to direct all internet-bound traffic to the

Deploy a NAT gateway in the public subnets. Modify the private subnet route table to direct all internet-bound traffic to the NAT gateway.

upvoted 4 times

ssa03 1 year, 3 months ago

Selected Answer: B

Correct Answer: B

upvoted 4 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Deploy a NAT gateway in the public subnets. Modify the private subnet route table to direct all internet-bound traffic to the NAT gateway.

upvoted 3 times

Deepakin96 1 year, 4 months ago

Selected Answer: B

NAT Gateway is safe

upvoted 3 times

Bmaster 1 year, 4 months ago

B is correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #550

Topic 1

A company is using AWS Key Management Service (AWS KMS) keys to encrypt AWS Lambda environment variables. A solutions architect needs to ensure that the required permissions are in place to decrypt and use the environment variables.

Which steps must the solutions architect take to implement the correct permissions? (Choose two.)

A. Add AWS KMS permissions in the Lambda resource policy.

B. Add AWS KMS permissions in the Lambda execution role. **Most Voted**

C. Add AWS KMS permissions in the Lambda function policy.

D. Allow the Lambda execution role in the AWS KMS key policy. **Most Voted**

E. Allow the Lambda resource policy in the AWS KMS key policy.

Correct Answer: BD

Community vote distribution

BD (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: BD

To decrypt environment variables encrypted with AWS KMS, Lambda needs to be granted permissions to call KMS APIs. This is done in two places:

The Lambda execution role needs kms:Decrypt and kms:GenerateDataKey permissions added. The execution role governs what AWS services the function code can access.

The KMS key policy needs to allow the Lambda execution role to have kms:Decrypt and kms:GenerateDataKey permissions for that specific key. This allows the execution role to use that particular key.

upvoted 7 times

wizcloudifa **Highly Voted** 7 months, 1 week ago

Selected Answer: BD

As per the principle of least privilege, granting permissions = role level

upvoted 5 times

TariqKipkemei **Most Recent** 1 year ago

Selected Answer: BD

Allow the Lambda execution role in the AWS KMS key policy then add AWS KMS permissions in the role.

upvoted 3 times

ssa03 1 year, 3 months ago

Selected Answer: BD

Correct Answer: BD

upvoted 3 times

Nirav1112 1 year, 4 months ago

its B & D

upvoted 2 times

mrsoa 1 year, 4 months ago

Selected Answer: BD

BD BD BD BD

upvoted 2 times

Deepakin96 1 year, 4 months ago

Selected Answer: BD

Its B and D

upvoted 2 times

Bmaster 1 year, 4 months ago

My choice is B,D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #551

Topic 1


A company has a financial application that produces reports. The reports average 50 KB in size and are stored in Amazon S3. The reports are frequently accessed during the first week after production and must be stored for several years. The reports must be retrievable within 6 hours.

Which solution meets these requirements MOST cost-effectively?

- A. Use S3 Standard. Use an S3 Lifecycle rule to transition the reports to S3 Glacier after 7 days. **Most Voted**
- B. Use S3 Standard. Use an S3 Lifecycle rule to transition the reports to S3 Standard-Infrequent Access (S3 Standard-IA) after 7 days.
- C. Use S3 Intelligent-Tiering. Configure S3 Intelligent-Tiering to transition the reports to S3 Standard-Infrequent Access (S3 Standard-IA) and S3 Glacier.
- D. Use S3 Standard. Use an S3 Lifecycle rule to transition the reports to S3 Glacier Deep Archive after 7 days.

Correct Answer: A

Community vote distribution



Comments

zjcorpuz **Highly Voted** 1 year, 10 months ago

Answer is A
 Amazon S3 Glacier:
 Expedited Retrieval: Provides access to data within 1-5 minutes.
 Standard Retrieval: Provides access to data within 3-5 hours.
 Bulk Retrieval: Provides access to data within 5-12 hours.
 Amazon S3 Glacier Deep Archive:
 Standard Retrieval: Provides access to data within 12 hours.
 Bulk Retrieval: Provides access to data within 48 hours.
 upvoted 26 times

oayoade **Highly Voted** 1 year, 9 months ago

Selected Answer: C

All the "...after 7 days" options are wrong.
 Before you transition objects to S3 Standard-IA or S3 One Zone-IA, you must store them for at least 30 days in Amazon S3
<https://docs.aws.amazon.com/AWSAmazonS3/latest/userguide/lifecycle-transition-general->

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html#:~:text=Minimum%20Days%20for%20Transition%20to%20S3%20Standard%2DIA%20or%20S3%20One%20Zone%2DIA>

upvoted 11 times

LeonSauveterre 6 months ago

Objects must be stored for at least 30 days before transitioning to S3 Standard-IA or S3 One Zone-IA.

Option A is to S3 Glacier. A is correct.

upvoted 2 times

MatAlves 8 months, 3 weeks ago

After reading the document from your link, it's clear that NONE of the restrictions apply to S3 standard, but to S3 Standard IA.

A is the correct answer.

upvoted 2 times

Hades2231 1 year, 9 months ago

This is worth noticing! Glad I came across your comment 1 day before my test.

upvoted 4 times

Marco_St 1 year, 5 months ago

so Could I ask is A or C for this question? I voted for A but it seems you had the same question in the exam and it was C? Thanks! I will attend the exam soon.

upvoted 2 times

franbarberan 1 year, 8 months ago

the 7 days limitation is only if you want to move from s3 standart to S3 Standard-IA or S3 One Zone-IA, if you move to s3 glacier dont have this limitation, correct answer is A

upvoted 14 times

TEC65 Most Recent 1 month, 2 weeks ago

Selected Answer: A

a is the answer

upvoted 1 times

rosanna 6 months ago

Selected Answer: A

there's a period before you can transition to any storage class (down the hierarchy) "except for glacier of all types" a minimum of 30 days.

However, when transitioning to Glacier directly there's no minimum period to wait so you can choose "A" safely

upvoted 1 times

1e22522 10 months ago

Selected Answer: A

Its A ya bunch of nerds

upvoted 3 times

Gape4 11 months, 2 weeks ago

Selected Answer: C

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 1 times

Linuslin 1 year, 1 month ago

Selected Answer: A

C is incorrect.

Unsupported lifecycle transitions

Amazon S3 does not support any of the following lifecycle transitions.

You can't transition from the following:

Any storage class to the S3 Standard storage class.

Any storage class to the Reduced Redundancy Storage (RRS) class.

The S3 Intelligent-Tiering storage class to the S3 Standard-IA storage class.

The S3 One Zone-IA storage class to the S3 Intelligent-Tiering, S3 Standard-IA, or S3 Glacier Instant Retrieval storage classes.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>

upvoted 4 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

BC are lifecycle with tiering and infrequent access which are not required here.

D is deep archive and can take hours to retrieve so it is not suitable

A is cheapest workable option

upvoted 4 times

Marco_St 1 year, 5 months ago

Selected Answer: A

frequent access pattern- Standard.

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: A

Not B - More expensive than A

Not C - Intelligent-Tiering moves only objects of at least 128 KB

Not D - Glacier Deep Archive takes more than 6 hours to retrieve

upvoted 4 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

Any option with S3 Intelligent-Tiering is out, this is only required when the access patterns are unknown.

From the question the access patterns are well known, enough to tie the frequently accessed reports to S3 standard and transition them to S3 glacier after 7days.

upvoted 4 times

iwannabeawsgod 1 year, 7 months ago

Selected Answer: A

its A for me

upvoted 3 times

Carlos_O 1 year, 8 months ago

Selected Answer: A

Tiene mas sentido

upvoted 2 times

sl2man 1 year, 8 months ago

Selected Answer: A

Option A

Amazon S3 Glacier Standard Retrieval: Provides access to data within 3-5 hours.

upvoted 4 times

Ramdi1 1 year, 8 months ago

Selected Answer: A

most cost effective has to be glacier so A

With C it is using intelligence tiering which is 30 days minimum from what I have read, I may be wrong on how I read that.

upvoted 2 times

tabbyDolly 1 year, 8 months ago

answer A

frequent access during the first week -> keeps data in s3 standard for 7 days
stored for several year and retrievable within 6 hours -> can be moved to s3 glacier for data archive purpose
upvoted 2 times

anikety123 1 year, 9 months ago

Selected Answer: A

Its A. Data cannot be transitioned from Intelligent Tiering to Standard IA
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html>
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #552

Topic 1

A company needs to optimize the cost of its Amazon EC2 instances. The company also needs to change the type and family of its EC2 instances every 2-3 months.

What should the company do to meet these requirements?

- A. Purchase Partial Upfront Reserved Instances for a 3-year term.
- B. Purchase a No Upfront Compute Savings Plan for a 1-year term. **Most Voted**
- C. Purchase All Upfront Reserved Instances for a 1-year term.
- D. Purchase an All Upfront EC2 Instance Savings Plan for a 1-year term.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Guru4Cloud **Highly Voted** 9 months, 3 weeks ago

Selected Answer: B

The key considerations are:

The company needs flexibility to change EC2 instance types and families every 2-3 months. This rules out Reserved Instances which lock you into an instance type and family for 1-3 years.

A Compute Savings Plan allows switching instance types and families freely within the term as needed. No Upfront is more flexible than All Upfront.

A 1-year term balances commitment and flexibility better than a 3-year term given the company's changing needs.

With No Upfront, the company only pays for usage monthly without an upfront payment. This optimizes cost.

upvoted 10 times

Kiki_Pass **Highly Voted** 10 months, 1 week ago

Selected Answer: B

"EC2 Instance Savings Plans give you the flexibility to change your usage between instances WITHIN a family in that region. "

<https://aws.amazon.com/savingsplans/compute-pricing/>

upvoted 5 times

TariqKipkemei **Most Recent** 6 months, 2 weeks ago

Selected Answer: B

Only Compute Savings Plan allows you to change instance family.
upvoted 4 times

avkya 10 months ago

Selected Answer: B

" needs to change the type and family of its EC2 instances". that means B I think.
upvoted 3 times

mrsoa 10 months, 1 week ago

Selected Answer: B

B is the right answer
upvoted 3 times

Bmaster 10 months, 1 week ago

B is correct..
'EC2 Instance Savings Plans' can't change 'family'.
upvoted 2 times

Josantru 10 months, 1 week ago

Correct B.
To change 'Family' always Compute saving plan, right?
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #553

Topic 1

A solutions architect needs to review a company's Amazon S3 buckets to discover personally identifiable information (PII). The company stores the PII data in the us-east-1 Region and us-west-2 Region.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Configure Amazon Macie in each Region. Create a job to analyze the data that is in Amazon S3. **Most Voted**
- B. Configure AWS Security Hub for all Regions. Create an AWS Config rule to analyze the data that is in Amazon S3.
- C. Configure Amazon Inspector to analyze the data that is in Amazon S3.
- D. Configure Amazon GuardDuty to analyze the data that is in Amazon S3.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: A

The key reasons are:

Amazon Macie is designed specifically for discovering and classifying sensitive data like PII in S3. This makes it the optimal service to use.

Macie can be enabled directly in the required Regions rather than enabling it across all Regions which is unnecessary. This minimizes overhead.

Macie can be set up to automatically scan the specified S3 buckets on a schedule. No need to create separate jobs.

Security Hub is for security monitoring across AWS accounts, not specific for PII discovery. More overhead than needed.

Inspector and GuardDuty are not built for PII discovery in S3 buckets. They provide broader security capabilities.

upvoted 6 times

awsgeek75 **Most Recent** 11 months ago

Selected Answer: A

PII = Macie

Security Hub: Organisation security and logging not for PII

Inspector: Infra vulnerability management

GuardDuty: Network protection

upvoted 4 times

TariqKipkemei 1 year ago

Selected Answer: A

Amazon Macie = PII

upvoted 2 times

mrsoa 1 year, 4 months ago

Selected Answer: A

AWS Macie = PII detection

upvoted 4 times

Deepakin96 1 year, 4 months ago

Selected Answer: A

Amazon Macie will identify all PII

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #554

Topic 1

A company's SAP application has a backend SQL Server database in an on-premises environment. The company wants to migrate its on-premises application and database server to AWS. The company needs an instance type that meets the high demands of its SAP database. On-premises performance data shows that both the SAP application and the database have high memory utilization.

Which solution will meet these requirements?

- A. Use the compute optimized instance family for the application. Use the memory optimized instance family for the database.
- B. Use the storage optimized instance family for both the application and the database.
- C. Use the memory optimized instance family for both the application and the database. **Most Voted**
- D. Use the high performance computing (HPC) optimized instance family for the application. Use the memory optimized instance family for the database.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Guru4Cloud** **Highly Voted** 9 months, 3 weeks ago**Selected Answer: C**

Since both the app and database have high memory needs, the memory optimized family like R5 instances meet those requirements well.

Using the same instance family simplifies management and operations, rather than mixing instance types.

Compute optimized instances may not provide enough memory for the SAP app's needs.

Storage optimized is overkill for the database's compute and memory needs.

HPC is overprovisioned for the SAP app.

upvoted 16 times

manOfThePeople **Highly Voted** 9 months, 1 week ago

High memory utilization = memory optimized.

C is the answer

upvoted 5 times

LeonSauveterre **Most Recent** 6 months ago

Selected Answer: C

Why D is incorrect: HPC family is geared toward compute-heavy workloads such as simulations or scientific calculations, not appropriate for a memory-intensive application like SAP.

upvoted 1 times

TariqKipkemei 6 months, 2 weeks ago

Selected Answer: C

Use the memory optimized instance family for both the application and the database

upvoted 3 times

mrsoa 10 months, 1 week ago

Selected Answer: C

I thynk its C

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #555

Topic 1

A company runs an application in a VPC with public and private subnets. The VPC extends across multiple Availability Zones. The application runs on Amazon EC2 instances in private subnets. The application uses an Amazon Simple Queue Service (Amazon SQS) queue.

A solutions architect needs to design a secure solution to establish a connection between the EC2 instances and the SQS queue.

Which solution will meet these requirements?

- A. Implement an interface VPC endpoint for Amazon SQS. Configure the endpoint to use the private subnets. Add to the endpoint a security group that has an inbound access rule that allows traffic from the EC2 instances that are in the private subnets. **Most Voted**
- B. Implement an interface VPC endpoint for Amazon SQS. Configure the endpoint to use the public subnets. Attach to the interface endpoint a VPC endpoint policy that allows access from the EC2 instances that are in the private subnets.
- C. Implement an interface VPC endpoint for Amazon SQS. Configure the endpoint to use the public subnets. Attach an Amazon SQS access policy to the interface VPC endpoint that allows requests from only a specified VPC endpoint.
- D. Implement a gateway endpoint for Amazon SQS. Add a NAT gateway to the private subnets. Attach an IAM role to the EC2 instances that allows access to the SQS queue.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: A

An interface VPC endpoint is a private way to connect to AWS services without having to expose your VPC to the public internet. This is the most secure way to connect to Amazon SQS from the private subnets. Configuring the endpoint to use the private subnets ensures that the traffic between the EC2 instances and the SQS queue is only within the VPC. This helps to protect the traffic from being intercepted by a malicious actor. Adding a security group to the endpoint that has an inbound access rule that allows traffic from the EC2 instances that are in the private subnets further restricts the traffic to only the authorized sources. This helps to prevent unauthorized access to the

SQS queue.

upvoted 11 times

Bmaster Highly Voted 1 year, 4 months ago

A is correct.

B,C: 'Configuring endpoints to use public subnets' --> Invalid

D: No Gateway Endpoint for SQS.

upvoted 6 times

awsgeek75 Most Recent 11 months ago

Selected Answer: A

BC are using public subnets so not useful for security

D uses gateway endpoint which is not useful to connect to SQS

A: <https://docs.aws.amazon.com/vpc/latest/privatelink/aws-services-privatelink-support.html>

upvoted 2 times

awsgeek75 10 months, 3 weeks ago

Sorry, the link is wrong for A. Please ignore it!

upvoted 2 times

ShawnTang 12 months ago

A seems the most suitable,

but security group can't add to the endpoint directly, right?

upvoted 2 times

TariqKipkemei 1 year ago

Selected Answer: A

Answer is A

upvoted 2 times

TariqKipkemei 1 year ago

Interface endpoints enable connectivity to services over AWS PrivateLink. It is a collection of one or more elastic network interfaces with a private IP address that serves as an entry point for traffic destined to a supported service.

Implement an interface VPC endpoint for Amazon SQS. Configure the endpoint to use the private subnets. Add to the endpoint a security group that has an inbound access rule that allows traffic from the EC2 instances that are in the private subnets.

upvoted 2 times

potomac 1 year, 1 month ago

A is correct

upvoted 2 times

mrsoa 1 year, 4 months ago

Selected Answer: A

I think its A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #556

Topic 1


A solutions architect is using an AWS CloudFormation template to deploy a three-tier web application. The web application consists of a web tier and an application tier that stores and retrieves user data in Amazon DynamoDB tables. The web and application tiers are hosted on Amazon EC2 instances, and the database tier is not publicly accessible. The application EC2 instances need to access the DynamoDB tables without exposing API credentials in the template.

What should the solutions architect do to meet these requirements?

- A. Create an IAM role to read the DynamoDB tables. Associate the role with the application instances by referencing an instance profile.
- B. Create an IAM role that has the required permissions to read and write from the DynamoDB tables. Add the role to the EC2 instance profile, and associate the instance profile with the application instances. **Most Voted**
- C. Use the parameter section in the AWS CloudFormation template to have the user input access and secret keys from an already-created IAM user that has the required permissions to read and write from the DynamoDB tables.
- D. Create an IAM user in the AWS CloudFormation template that has the required permissions to read and write from the DynamoDB tables. Use the GetAtt function to retrieve the access and secret keys, and pass them to the application instances through the user data.

Correct Answer: B

Community vote distribution

**Comments**

upliftinghut 10 months, 2 weeks ago

Selected Answer: B

best practice is using IAM role for database access. From app to DB => need both read & write, only B meets these 2
upvoted 3 times

pentium75 11 months, 1 week ago

Selected Answer: B

Application "stores and retrieves" data in DynamoDB while A grants only access "to read".
upvoted 3 times

Nisarg2121 1 year, 1 month ago

Selected Answer: B

B is correct, A total wrong because "read the DynamoDB tables", so what about write in database.

upvoted 4 times

darekw 1 year, 3 months ago

question says: ...application tier stores and retrieves user data in Amazon DynamoDB tables... so it needs read and write access

A) is only read access

B) seems to be the right answer

upvoted 2 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

Option B is the correct approach to meet the requirements:

Create an IAM role with permissions to access DynamoDB

Add the IAM role to an EC2 Instance Profile

Associate the Instance Profile with the application EC2 instances

This allows the instances to assume the IAM role to obtain temporary credentials to access DynamoDB.

upvoted 3 times

anibinaadi 1 year, 3 months ago

Explanation. Both A and B seems suitable. But Option A is incorrect because it says "Associate the role with the application instances by referencing an instance profile". Which just only a Part of the solution.

In API/AWS CLI following steps are required to complete the Role-> instance profile association-> to instance.

1. Create an IAM Role

2. add-role-to-instance-profile (aws iam add-role-to-instance-profile --role-name S3Access --instance-profile-name Webserver)

3. associate-iam-instance-profile (aws ec2 associate-iam-instance-profile --instance-id i-123456789abcde123 --iam-instance-profile Name=admin-role)

hence Option B is correct.

upvoted 3 times

DannyKang5649 1 year, 3 months ago

Selected Answer: B

Why "No read and write" ? The question clearly states that application tier STORE and RETRIEVE the data from DynamoDB.

Which means write and read... I think answer should be B

upvoted 3 times

xyb 1 year, 4 months ago

Selected Answer: B

<https://www.examttopics.com/discussions/amazon/view/80755-exam-aws-certified-solutions-architect-associate-saa-c02/>

upvoted 2 times

Ale1973 1 year, 4 months ago

Selected Answer: B

My rationl: Option A is wrong because the scenario says "stores and retrieves user data in Amazon DynamoDB tables", STORES and RETRIVE, if you set a role to READ, you can write on DinamoDB database

upvoted 2 times

mrsoa 1 year, 4 months ago

Selected Answer: A

AAAAAAAAA

upvoted 2 times

pentium75 11 months, 1 week ago

No because it grants only read access

upvoted 3 times

kangho 1 year, 4 months ago

Selected Answer: A

A is correct
upvoted 1 times

pentium75 11 months, 1 week ago

No because it grants only read access
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #557

Topic 1

A solutions architect manages an analytics application. The application stores large amounts of semistructured data in an Amazon S3 bucket. The solutions architect wants to use parallel data processing to process the data more quickly. The solutions architect also wants to use information that is stored in an Amazon Redshift database to enrich the data.

Which solution will meet these requirements?

- A. Use Amazon Athena to process the S3 data. Use AWS Glue with the Amazon Redshift data to enrich the S3 data.
- B. Use Amazon EMR to process the S3 data. Use Amazon EMR with the Amazon Redshift data to enrich the S3 data.**
Most Voted
- C. Use Amazon EMR to process the S3 data. Use Amazon Kinesis Data Streams to move the S3 data into Amazon Redshift so that the data can be enriched.
- D. Use AWS Glue to process the S3 data. Use AWS Lake Formation with the Amazon Redshift data to enrich the S3 data.

Correct Answer: B*Community vote distribution***Comments****Guru4Cloud** **Highly Voted** 1 year, 3 months ago**Selected Answer: B**

Option B is the correct solution that meets the requirements:

Use Amazon EMR to process the semi-structured data in Amazon S3. EMR provides a managed Hadoop framework optimized for processing large datasets in S3.

EMR supports parallel data processing across multiple nodes to speed up the processing.

EMR can integrate directly with Amazon Redshift using the EMR-Redshift integration. This allows querying the Redshift data from EMR and joining it with the S3 data.

This enables enriching the semi-structured S3 data with the information stored in Redshift

upvoted 18 times

zjcorpuz **Highly Voted** 1 year, 4 months ago

By combining AWS Glue and Amazon Redshift, you can process the semistructured data in parallel using Glue ETL jobs and then store the processed and enriched data in a structured format in Amazon Redshift. This approach allows you to perform complex analytics efficiently and at scale.

complex analytics efficiently and at scale.

upvoted 9 times

upliftinghut Most Recent 10 months, 2 weeks ago

Selected Answer: B

D: not relevant, data is semistructured and Glue is more batch than stream data

A: not correct, Athena is for querying data

B & C look ok but C is out => redundant with Kinesis data stream; EMR already processed data as input into Redshift for parallel processing

Only B is most logical

upvoted 4 times

awsgeek75 10 months, 3 weeks ago

Selected Answer: B

Key requirement: parallel data processing

parallel data processing is EMR (Kind of Apache Hadoop) so it only leave B and C

C is Kinesis to Redshift which is pointless logic here

B EMR for S3 and EMR for Redshift gives maximum parallel processing here

upvoted 3 times

pentium75 11 months, 1 week ago

Selected Answer: B

A has a pitfall, "use Amazon Athena to PROCESS the data". With Athena you can query, not process, data.

C is wrong because Kinesis has no place here.

D is wrong because it does not process the Redshift data, and Glue does ETL, not analyze

Thus it's B. EMR can use semi-structured data from from S3 and structured data from Redshift and is ideal for "parallel data processing" of "large amounts" of data.

upvoted 8 times

aws94 12 months ago

Selected Answer: B

large amount of data + parallel data processing = EMR

upvoted 3 times

[Removed] 12 months ago

Selected Answer: A

Amazon Athena is an interactive query service that makes it easy to analyze data directly in Amazon Simple Storage Service (Amazon S3) using standard SQL.

upvoted 1 times

pentium75 11 months, 1 week ago

Y, but A says "process", not "query" data with Athena.

upvoted 2 times

SHAAHIBHUSHANAWS 1 year ago

Selected Answer: D

Glue use apache pyspark cluster for parallel processing. EMR or Glue are possible options. Glue is serverless so better use this plus pyspark is in memory parallel processing.

upvoted 1 times

aragornfsm 1 year ago

i think a is correct

semistructured data ==> Athena

upvoted 1 times

pentium75 11 months, 1 week ago

"Hadoop [as used by EMR] helps you turn petabytes of un-structured or semi-structured data into useful insights"

<https://aws.amazon.com/emr/features/hadoop/>
upvoted 2 times

riyasara 1 year ago

Athena is not designed for parallel data processing. So it's B
upvoted 3 times

TariqKipkemei 1 year ago

Selected Answer: A

Answer is A
upvoted 1 times

TariqKipkemei 1 year ago

Selected Answer: B

From this documentation looks like EMR cannot interface with S3.
<https://aws.amazon.com/emr/>

I will settle with option A.
upvoted 2 times

pentium75 11 months, 1 week ago

Of course EMR can access S3

<https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-plan-file-systems.html>
upvoted 2 times

bogobob 1 year ago

Selected Answer: B

For those answering A, AWS Glue can directly query S3, it can't use Athena as a source of data. The questions say the Redshift data should be user to "enrich" which means thats the redshift data needs to be "added" to the s3 data. A doesn't allow that.
upvoted 2 times

hungta 1 year ago

Selected Answer: B

Choose option B.
Option A is not correct. Amazon Athena is suitable for querying data directly from S3 using SQL and allows parallel processing of S3 data.
AWS Glue can be used for data preparation and enrichment but might not directly integrate with Amazon Redshift for enrichment.
upvoted 2 times

potomac 1 year, 1 month ago

Selected Answer: A

Athena and Redshift both do SQL query
upvoted 1 times

Sab123 1 year, 2 months ago

Selected Answer: A

semi-structure supported by Athena not by EMR
upvoted 4 times

pentium75 11 months, 1 week ago

"Hadoop helps you turn petabytes of un-structured or semi-structured data into useful insights about your applications or users."

https://aws.amazon.com/emr/features/hadoop/?nc1=h_ls
upvoted 2 times

Wade770 1 year, 8 months ago

JKevin / 8 1 year, 2 months ago

Selected Answer: A

athena for s3

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #558

Topic 1

A company has two VPCs that are located in the us-west-2 Region within the same AWS account. The company needs to allow network traffic between these VPCs. Approximately 500 GB of data transfer will occur between the VPCs each month.

What is the MOST cost-effective solution to connect these VPCs?

- A. Implement AWS Transit Gateway to connect the VPCs. Update the route tables of each VPC to use the transit gateway for inter-VPC communication.
- B. Implement an AWS Site-to-Site VPN tunnel between the VPCs. Update the route tables of each VPC to use the VPN tunnel for inter-VPC communication.
- C. Set up a VPC peering connection between the VPCs. Update the route tables of each VPC to use the VPC peering connection for inter-VPC communication. **Most Voted**
- D. Set up a 1 GB AWS Direct Connect connection between the VPCs. Update the route tables of each VPC to use the Direct Connect connection for inter-VPC communication.

Correct Answer: C

Community vote distribution

C (100%)

Comments

luiscc **Highly Voted** 10 months, 1 week ago

Selected Answer: C

C is the correct answer.

VPC peering is the most cost-effective way to connect two VPCs within the same region and AWS account. There are no additional charges for VPC peering beyond standard data transfer rates.

Transit Gateway and VPN add additional hourly and data processing charges that are not necessary for simple VPC peering.

Direct Connect provides dedicated network connectivity, but is overkill for the relatively low inter-VPC data transfer needs described here. It has high fixed costs plus data transfer rates.

For occasional inter-VPC communication of moderate data volumes within the same region and account, VPC peering is the most cost-effective solution. It provides simple private connectivity without transfer charges or network appliances.

upvoted 5 times

Guru4Cloud Highly Voted 9 months, 3 weeks ago

Selected Answer: C

The key reasons are:

VPC peering provides private connectivity between VPCs without using public IP space.

Data transferred between peered VPCs is free as long as they are in the same region.

500 GB/month inter-VPC data transfer fits within peering free tier.

Transit Gateway (Option A) incurs hourly charges plus data transfer fees. More costly than peering.

Site-to-Site VPN (Option B) incurs hourly charges and data transfer fees. More expensive than peering.

Direct Connect (Option D) has high hourly charges and would be overkill for this use case.

upvoted 5 times

TariqKipkemei Most Recent 6 months, 2 weeks ago

Selected Answer: C

A VPC peering connection is a networking connection between two VPCs that enables you to route traffic between them using private IPv4 addresses or IPv6 addresses. Instances in either VPC can communicate with each other as if they are within the same network. You can create a VPC peering connection between your own VPCs, or with a VPC in another AWS account. The VPCs can be in different Regions (also known as an inter-Region VPC peering connection).

<https://docs.aws.amazon.com/vpc/latest/peering/what-is-vpc-peering.html#:~:text=A-,VPC%20peering,-connection%20is%20a>

upvoted 3 times

BrijMohan08 9 months, 2 weeks ago

Selected Answer: C

Transit Gateway network peering.

VPC Peering to peer 2 or more VPC in the same region.

upvoted 4 times

mrsoa 10 months, 1 week ago

Selected Answer: C

VPC peering is the most cost-effective solution

upvoted 2 times

Deepakin96 10 months, 1 week ago

Selected Answer: C

Communicating with two VPC in same account = VPC Peering

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #559

Topic 1

A company hosts multiple applications on AWS for different product lines. The applications use different compute resources, including Amazon EC2 instances and Application Load Balancers. The applications run in different AWS accounts under the same organization in AWS Organizations across multiple AWS Regions. Teams for each product line have tagged each compute resource in the individual accounts.

The company wants more details about the cost for each product line from the consolidated billing feature in Organizations.

Which combination of steps will meet these requirements? (Choose two.)

A. Select a specific AWS generated tag in the AWS Billing console.

B. Select a specific user-defined tag in the AWS Billing console. **Most Voted**

C. Select a specific user-defined tag in the AWS Resource Groups console.

D. Activate the selected tag from each AWS account.

E. Activate the selected tag from the Organizations management account. **Most Voted**

Correct Answer: BE

Community vote distribution

BE (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: BE

The reasons are:

User-defined tags were created by each product team to identify resources. Selecting the relevant tag in the Billing console will group costs.

The tag must be activated from the Organizations management account to consolidate billing across all accounts.

AWS generated tags are predefined by AWS and won't align to product lines.

Resource Groups (Option C) helps manage resources but not billing.

Activating the tag from each account (Option D) is not needed since Organizations centralizes billing.

upvoted 11 times

adiorsdior999 7 months, 1 week ago

Resource Groups helps grouping, editing, searching tag for resources, and do not use for viewing cost details
upvoted 2 times

potomac Most Recent 1 year, 7 months ago

Selected Answer: BE

Your user-defined cost allocation tags represent the tag key, which you activate in the Billing console.
upvoted 2 times

mrsoa 1 year, 10 months ago

Selected Answer: BE

BE BE BE BE
upvoted 3 times

Kiki_Pass 1 year, 10 months ago

Selected Answer: BE

"Only a management account in an organization and single accounts that aren't members of an organization have access to the cost allocation tags manager in the Billing and Cost Management console."
<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/custom-tags.html>
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #560


Topic 1

A company's solutions architect is designing an AWS multi-account solution that uses AWS Organizations. The solutions architect has organized the company's accounts into organizational units (OUs).

The solutions architect needs a solution that will identify any changes to the OU hierarchy. The solution also needs to notify the company's operations team of any changes.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Provision the AWS accounts by using AWS Control Tower. Use account drift notifications to identify the changes to the OU hierarchy. **Most Voted**
- B. Provision the AWS accounts by using AWS Control Tower. Use AWS Config aggregated rules to identify the changes to the OU hierarchy.
- C. Use AWS Service Catalog to create accounts in Organizations. Use an AWS CloudTrail organization trail to identify the changes to the OU hierarchy.
- D. Use AWS CloudFormation templates to create accounts in Organizations. Use the drift detection operation on a stack to identify the changes to the OU hierarchy.

Correct Answer: A*Community vote distribution***Comments****Guru4Cloud** **Highly Voted** 1 year, 9 months ago**Selected Answer: A**

The key advantages you highlight of Control Tower are convincing:

Fully managed service simplifies multi-account setup.
Built-in account drift notifications detect OU changes automatically.
More scalable and less complex than Config rules or CloudTrail.
Better security and compliance guardrails than custom options.
Lower operational overhead compared to other solution

upvoted 12 times

Bmaster Highly Voted 1 year, 10 months ago

A is correct.

<https://docs.aws.amazon.com/controltower/latest/userguide/what-is-control-tower.html>

<https://docs.aws.amazon.com/controltower/latest/userguide/prevention-and-notification.html>

upvoted 8 times

1166ae3 Most Recent 11 months, 1 week ago

Selected Answer: C

Create Accounts using AWS Service Catalog:

Utilize AWS Service Catalog to provision AWS accounts within AWS Organizations. This ensures standardized account creation and management.

Enable AWS CloudTrail Organization Trail:

Set up an AWS CloudTrail organization trail that records all API calls across all accounts in the organization.

This trail will capture changes to the OU hierarchy, including any modifications to organizational units.

upvoted 2 times

chickenmf 1 year, 2 months ago

Selected Answer: B

AWS Config helps you maintain a detailed inventory of your resources and their configurations, track changes over time, and ensure compliance with your organization's policies and industry regulations.

upvoted 4 times

chickenmf 1 year, 2 months ago

Furthermore, AWS Config Aggregated Rules are a feature within AWS Config that enables you to evaluate compliance with desired configurations or compliance standards across multiple AWS accounts and regions. They are particularly useful in scenarios where you want to enforce consistent rules and compliance checks across an entire organization with multiple AWS accounts.

upvoted 1 times

chickenmf 1 year, 2 months ago

NVM - This is such a stupid question lol changing my answer to A due to the following:

Account drift notifications in AWS are a feature provided by AWS Control Tower. These notifications help organizations identify and respond to changes made to an AWS account that deviate from the established baseline configuration created during the initial setup by AWS Control Tower. Drift refers to any configuration changes that have been made to an AWS account after it was provisioned by Control Tower.

upvoted 5 times

Avyay 1 year, 3 months ago

This was in my exam today..I selected Answer A

upvoted 3 times

chickenmf 1 year, 2 months ago

what percentage of all these questions would you say were in the exam?

upvoted 2 times

wizcloudifa 1 year, 1 month ago

I read in one of the earlier questions, its around 75%, someone who gave the exam said so

upvoted 2 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

A

<https://docs.aws.amazon.com/controltower/latest/userguide/drift.html>

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: A

AWS Control Tower provides passive and active methods of drift monitoring protection for preventive controls.
upvoted 2 times

darekw 1 year, 9 months ago

<https://docs.aws.amazon.com/controltower/latest/userguide/prevention-and-notification.html>
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #561

Topic 1

A company's website handles millions of requests each day, and the number of requests continues to increase. A solutions architect needs to improve the response time of the web application. The solutions architect determines that the application needs to decrease latency when retrieving product details from the Amazon DynamoDB table.

Which solution will meet these requirements with the LEAST amount of operational overhead?

- A. Set up a DynamoDB Accelerator (DAX) cluster. Route all read requests through DAX. **Most Voted**
- B. Set up Amazon ElastiCache for Redis between the DynamoDB table and the web application. Route all read requests through Redis.
- C. Set up Amazon ElastiCache for Memcached between the DynamoDB table and the web application. Route all read requests through Memcached.
- D. Set up Amazon DynamoDB Streams on the table, and have AWS Lambda read from the table and populate Amazon ElastiCache. Route all read requests through ElastiCache.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**mrsoa** **Highly Voted** 1 year, 4 months ago**Selected Answer: A**

A , because B,C and D contains Elasicache which required a heavy code changes, so more operational overhead
upvoted 9 times

TariqKipkemei **Highly Voted** 1 year ago**Selected Answer: A**

decrease latency when retrieving product details from the Amazon DynamoDB = Amazon DynamoDB Accelerator (DAX)
upvoted 5 times

sandordini **Most Recent** 7 months, 3 weeks ago**Selected Answer: A**

DAX to reduce latency
upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: A

The key reasons:

DAX provides a DynamoDB-compatible caching layer to reduce read latency. It is purpose-built for accelerating DynamoDB workloads.

Using DAX requires minimal application changes - only read requests are routed through it.

DAX handles caching logic automatically without needing complex integration code.

ElastiCache Redis/Memcached (Options B/C) require more integration work to sync DynamoDB data.

Using Lambda and Streams to populate ElastiCache (Option D) is a complex event-driven approach requiring ongoing maintenance.

DAX plugs in seamlessly to accelerate DynamoDB with very little operational overhead

upvoted 3 times

Deepakin96 1 year, 4 months ago

Selected Answer: A

DynamoDB = DAX

upvoted 3 times

Bmaster 1 year, 4 months ago

only A

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #562

Topic 1

A solutions architect needs to ensure that API calls to Amazon DynamoDB from Amazon EC2 instances in a VPC do not travel across the internet.

Which combination of steps should the solutions architect take to meet this requirement? (Choose two.)

- A. Create a route table entry for the endpoint. **Most Voted**
- B. Create a gateway endpoint for DynamoDB. **Most Voted**
- C. Create an interface endpoint for Amazon EC2.
- D. Create an elastic network interface for the endpoint in each of the subnets of the VPC.
- E. Create a security group entry in the endpoint's security group to provide access.

Correct Answer: AB*Community vote distribution***Comments****ukivanlampi** **Highly Voted** 1 year, 10 months ago**Selected Answer: AB**

<https://docs.aws.amazon.com/vpc/latest/privatelink/vpc-endpoints-ddb.html>
upvoted 12 times

Guru4Cloud **Highly Voted** 1 year, 9 months ago**Selected Answer: BE**

The reasons are:

A gateway endpoint for DynamoDB enables private connectivity between DynamoDB and the VPC. This allows EC2 instances to access DynamoDB APIs without traversing the internet.

A security group entry is needed to allow the EC2 instances access to the DynamoDB endpoint over the VPC.

An interface endpoint is used for services like S3 and Systems Manager, not DynamoDB.

Route table entries route traffic within a VPC but do not affect external connectivity.

Elastic network interfaces are not needed for gateway endpoints.

upvoted 10 times

unbendable 1 year, 7 months ago

"The outbound rules for the security group for instances that access DynamoDB through the gateway endpoint must allow traffic to DynamoDB", <https://docs.aws.amazon.com/vpc/latest/privatelink/vpc-endpoints-ddb.html>

The option however is talking about the security group of the endpoint

upvoted 3 times

JoeTromundo Most Recent 8 months, 2 weeks ago

Selected Answer: AB

A & B are correct

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html>

E is incorrect. There's no need for security group.

From the URL above:

"Once the VPC subnet's gateway endpoint has been granted access to DynamoDB, any AWS account with access to that subnet can use DynamoDB."

upvoted 2 times

a7md0 11 months, 1 week ago

Selected Answer: AB

Creating the gateway endpoint and edit the route table is enough, there are no security group involved

upvoted 2 times

osmk 1 year, 4 months ago

AB <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html>

upvoted 3 times

upliftinghut 1 year, 4 months ago

Selected Answer: AB

C & D are both not relevant. D looks ok but DynamoDB doesn't go with security group, it only allows route table for VPC endpoint. Link here: <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html>

upvoted 2 times

upliftinghut 1 year, 4 months ago

Sorry, E not D. E looks ok but DynamoDB doesn't go with security group, it only allows route table for VPC endpoint. Link here: <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html>

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: AB

DynamoDB can only be connected via Gateway endpoint (just like S3)

route table for connecting the VPC for the endpoint

So do B then A

C: interface endpoint for EC2 to what?

D: ENI not applicable here for VPC

E: Incomplete option as to access to what?

upvoted 4 times

theonlyhero 1 year, 4 months ago

go through this video it will show the answer is AB

<https://www.youtube.com/watch?v=8FTnyhklEvU>

upvoted 5 times

pentium75 1 year, 5 months ago

Selected Answer: AB

Gateway Endpoint does not have an ENI, thus it has no security group. Instances have security groups and those must allow access to DynamoDB.

upvoted 6 times

aws94 1 year, 5 months ago

Selected Answer: BE

A. Create a route table entry for the endpoint: This is not necessary, as the gateway endpoint itself automatically creates the required route table entries.

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: AB

Create a gateway endpoint for DynamoDB then create a route table entry for the endpoint

upvoted 3 times

EdenWang 1 year, 6 months ago

Selected Answer: BE

refer to question 555

upvoted 2 times

cciesam 1 year, 6 months ago

Selected Answer: AB

<https://docs.aws.amazon.com/vpc/latest/privatelink/gateway-endpoints.html#vpc-endpoints-routing>

Traffic from your VPC to Amazon S3 or DynamoDB is routed to the gateway endpoint. Each subnet route table must have a route that sends traffic destined for the service to the gateway endpoint using the prefix list for the service.

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: AB

You can access Amazon DynamoDB from your VPC using gateway VPC endpoints. After you create the gateway endpoint, you can add it as a target in your route table for traffic destined from your VPC to DynamoDB.

upvoted 3 times

danielmakita 1 year, 7 months ago

It is A and B. Not E because security group does not span VPCs.

upvoted 3 times

iwannabeawsgod 1 year, 7 months ago

Selected Answer: AB

A and B for sure

upvoted 4 times

loveaws 1 year, 8 months ago

B and D.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #563

Topic 1


A company runs its applications on both Amazon Elastic Kubernetes Service (Amazon EKS) clusters and on-premises Kubernetes clusters. The company wants to view all clusters and workloads from a central location.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon CloudWatch Container Insights to collect and group the cluster information.
- B. Use Amazon EKS Connector to register and connect all Kubernetes clusters. **Most Voted**
- C. Use AWS Systems Manager to collect and view the cluster information.
- D. Use Amazon EKS Anywhere as the primary cluster to view the other clusters with native Kubernetes commands.

Correct Answer: B

Community vote distribution



Comments

ErnShm **Highly Voted** 1 year, 9 months ago

B

You can use Amazon EKS Connector to register and connect any conformant Kubernetes cluster to AWS and visualize it in the Amazon EKS console. After a cluster is connected, you can see the status, configuration, and workloads for that cluster in the Amazon EKS console. You can use this feature to view connected clusters in Amazon EKS console, but you can't manage them. The Amazon EKS Connector requires an agent that is an open source project on Github. For additional technical content, including frequently asked questions and troubleshooting, see [Troubleshooting issues in Amazon EKS Connector](#)

The Amazon EKS Connector can connect the following types of Kubernetes clusters to Amazon EKS.

On-premises Kubernetes clusters

Self-managed clusters that are running on Amazon EC2

Managed clusters from other cloud providers

upvoted 7 times

mrsoa **Highly Voted** 1 year, 10 months ago

seems B

<https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html>

upvoted 5 times

Danilus Most Recent 7 months ago

Selected Answer: B

KEY-he company wants to view all clusters and workloads from a central location.

the answer is not D because amazon EKS anywhere allow you to run cluster of kubernetes in on-premises or in other clouds
the answer is B EKS connector helps u to connect kubernetes clusters on premises or other cloud and manage from a single
ubication

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

<https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html>

"You can use Amazon EKS Connector to register and connect any conformant Kubernetes cluster to AWS and visualize it in the Amazon EKS console. "

B is the right product for this.

upvoted 4 times

pentium75 1 year, 5 months ago

Selected Answer: B

EKS Connector -> 'view clusters and workloads' as requested

EKS Anywhere -> create and manage on-premises EKS clusters

upvoted 4 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

B
EKS connector helps to integrate multiple cluster with EKS console. EKS anywhere is Kubernetes Ditro cluster to be deployed on-
prem. It is not for integrating with other cluster.

upvoted 1 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

View all clusters and workloads (incl on-prem) from a central location = Amazon EKS Connector

Create and operate Kubernetes clusters on your own infrastructure = Amazon EKS Anywhere

<https://aws.amazon.com/eks/eks-anywhere/#:~:text=Amazon-,EKS%20Anywhere,-lets%20you%20create>

<https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html#:~:text=You%20can%20use-,Amazon%20EKS%20Connector,-to%20register%20and>

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: B

It is B

upvoted 2 times

thainguyensunya 1 year, 9 months ago

Selected Answer: B

Definitely B.

"You can use Amazon EKS Connector to register and connect any conformant Kubernetes cluster to AWS and visualize it in the Amazon EKS console. After a cluster is connected, you can see the status, configuration, and workloads for that cluster in the Amazon EKS console. "

<https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html>

upvoted 3 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

The key reasons:

EKS Connector allows registering external Kubernetes clusters (on-premises and otherwise) with Amazon EKS. This provides a unified view and management of all clusters within the EKS console. EKS Connector handles keeping resources in sync across connected clusters. This centralized approach minimizes operational overhead compared to using separate tools. CloudWatch Container Insights (Option A) only provides metrics and logs, not cluster management. Systems Manager (Option C) is more general purpose and does not natively integrate with EKS. EKS Anywhere (Option D) would not provide a single pane of glass for external clusters.

upvoted 4 times

RealMarcus 1 year, 10 months ago

Amazon EKS Connector enables you to create and manage a centralized view of all your Kubernetes clusters, regardless of whether they are Amazon EKS clusters or on-premises Kubernetes clusters. It allows you to register these clusters with your Amazon EKS control plane, providing a unified management interface for all clusters.

upvoted 2 times

avkya 1 year, 10 months ago

Selected Answer: B

You can use Amazon EKS Connector to register and connect any conformant Kubernetes cluster to AWS and visualize it in the Amazon EKS console. After a cluster is connected, you can see the status, configuration, and workloads for that cluster in the Amazon EKS console. You can use this feature to view connected clusters in Amazon EKS console, but you can't manage them

upvoted 1 times

ukivanlamipi 1 year, 10 months ago

Selected Answer: D

only D can connect to on-perm

upvoted 1 times

pentium75 1 year, 5 months ago

No.

"The Amazon EKS Connector can connect the following types of Kubernetes clusters to Amazon EKS.

On-premises Kubernetes clusters"

<https://aws.amazon.com/de/eks/eks-anywhere/>

upvoted 2 times

pentium75 1 year, 5 months ago

Wrong link, statement is from <https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html>

upvoted 2 times

Bmaster 1 year, 10 months ago

Only B

<https://docs.aws.amazon.com/eks/latest/userguide/eks-connector.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #564

Topic 1

A company is building an ecommerce application and needs to store sensitive customer information. The company needs to give customers the ability to complete purchase transactions on the website. The company also needs to ensure that sensitive customer data is protected, even from database administrators.

Which solution meets these requirements?

- A. Store sensitive data in an Amazon Elastic Block Store (Amazon EBS) volume. Use EBS encryption to encrypt the data. Use an IAM instance role to restrict access.
- B. Store sensitive data in Amazon RDS for MySQL. Use AWS Key Management Service (AWS KMS) client-side encryption to encrypt the data. **Most Voted**
- C. Store sensitive data in Amazon S3. Use AWS Key Management Service (AWS KMS) server-side encryption to encrypt the data. Use S3 bucket policies to restrict access.
- D. Store sensitive data in Amazon FSx for Windows Server. Mount the file share on application servers. Use Windows file permissions to restrict access.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: B

The key reasons are:

RDS MySQL provides a fully managed database service well suited for an ecommerce application. AWS KMS client-side encryption allows encrypting sensitive data before it hits the database. The data remains encrypted at rest.

This protects sensitive customer data from database admins and privileged users.

EBS encryption (Option A) protects data at rest but not in use. IAM roles don't prevent admin access.

S3 (Option C) encrypts data at rest on the server side. Bucket policies don't restrict admin access.

FSx file permissions (Option D) don't prevent admin access to unencrypted data.

upvoted 12 times

pentium75 **Highly Voted** 11 months, 1 week ago

Selected Answer: B

A, C and D would allow the administrator of the storage to access the data. Besides, it is data about "purchase transactions" which is usually stored in a transactional database (such as RDS for MySQL), not in a file or object storage.

upvoted 8 times

sidharthwader 7 months ago

Good thought purchase transactions I had missed that part

upvoted 2 times

SHAAHIBHUSHANAWS Most Recent 1 year ago

B

I want to go with B as question is for database administrator. Also client key encryption is possible in code and KMS can be used for encryption but not using KMS keys. Encrypted data available in DB is of no use to DB admin.

upvoted 2 times

riyasara 1 year ago

Answer is option C. option B is not ideal because Amazon RDS for MySQL is a relational database service that is optimized for structured data, not for storing sensitive customer information. Moreover, by using client-side encryption with AWS KMS, you need to encrypt and decrypt the data in your application code, which increases the risk of exposing your data in transit or at rest. You also need to manage the encryption keys yourself, which adds complexity and overhead to your application.

upvoted 2 times

awsgeek75 11 months ago

eCommerce data and transaction data are ideal for RDS which, when encrypted, is secure even from the DBA.

upvoted 2 times

pentium75 11 months, 1 week ago

"optimized for structured data, not for storing sensitive customer information" ... Data related to "purchase transactions" is usually structured; that it contains "sensitive customer information" doesn't change the structured nature.

upvoted 4 times

wsdasdasdqwdaw 1 year, 1 month ago

I would go for B, because RDS (database admins), but I would like to see as well encryption at rest as well, not only in transit.

upvoted 2 times

mrsoa 1 year, 4 months ago

Selected Answer: B

Using client-side encryption we can protect specific fields and guarantee only decryption if the client has access to an API key, we can protect specific fields even from database admins

upvoted 3 times

D10SJoker 1 year, 4 months ago

Selected Answer: B

For me it's B because of "client-side encryption to encrypt the data"

upvoted 2 times

h8er 1 year, 4 months ago

keyword - database administrators

upvoted 5 times

Kiki_Pass 1 year, 4 months ago

Selected Answer: B

"even from database administrators" -> "Client Side encryption"

upvoted 4 times

Bmaster 1 year, 4 months ago

My choice is B
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #565

Topic 1

A company has an on-premises MySQL database that handles transactional data. The company is migrating the database to the AWS Cloud. The migrated database must maintain compatibility with the company's applications that use the database. The migrated database also must scale automatically during periods of increased demand.

Which migration solution will meet these requirements?

- A. Use native MySQL tools to migrate the database to Amazon RDS for MySQL. Configure elastic storage scaling.
- B. Migrate the database to Amazon Redshift by using the mysqldump utility. Turn on Auto Scaling for the Amazon Redshift cluster.
- C. Use AWS Database Migration Service (AWS DMS) to migrate the database to Amazon Aurora. Turn on Aurora Auto Scaling. **Most Voted**
- D. Use AWS Database Migration Service (AWS DMS) to migrate the database to Amazon DynamoDB. Configure an Auto Scaling policy.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Guru4Cloud** **Highly Voted** 1 year, 3 months ago**Selected Answer: C**

The key reasons are:

DMS provides an easy migration path from MySQL to Aurora while minimizing downtime. Aurora is a MySQL-compatible relational database service that will maintain compatibility with the company's applications. Aurora Auto Scaling allows the database to automatically scale up and down based on demand to handle increased workloads. RDS MySQL (Option A) does not scale as well as the Aurora architecture. Redshift (Option B) is for analytics, not transactional data, and may not be compatible. DynamoDB (Option D) is a NoSQL datastore and lacks MySQL compatibility.

upvoted 9 times

awsgeek75 **Most Recent** 11 months ago**Selected Answer: C**

A is wrong as you cannot use native MySQL tools for migration. Happy to be corrected though!

B Redshift is not compatible with MySQL

D is DynamoDB

C Aurora MySQL is compatible and supports auto scaling

upvoted 3 times

TariqKipkemei 1 year ago

Selected Answer: C

on-premises MySQL database, transactional data, maintain compatibility, scale automatically = Amazon Aurora
migrating the database to the AWS Cloud = AWS Database Migration Service

upvoted 2 times

potomac 1 year, 1 month ago

Selected Answer: C

Aurora is a MySQL-compatible relational database service

upvoted 2 times

mrsoa 1 year, 4 months ago

Selected Answer: C

Aurora is better in autoscaling then RDS

upvoted 2 times

Bmaster 1 year, 4 months ago

C is correct

A is incorrect. RDS for MySQL does not scale automatically during periods of increased demand.

B is incorrect. Redshift is used for data sharing purposes.

D is incorrect. you must change application codes.

upvoted 2 times

Eminenza22 1 year, 4 months ago

Amazon RDS now supports Storage Auto Scaling

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #566

Topic 1

A company runs multiple Amazon EC2 Linux instances in a VPC across two Availability Zones. The instances host applications that use a hierarchical directory structure. The applications need to read and write rapidly and concurrently to shared storage.

What should a solutions architect do to meet these requirements?

A. Create an Amazon S3 bucket. Allow access from all the EC2 instances in the VPC.

B. Create an Amazon Elastic File System (Amazon EFS) file system. Mount the EFS file system from each EC2 instance.


Most Voted

C. Create a file system on a Provisioned IOPS SSD (io2) Amazon Elastic Block Store (Amazon EBS) volume. Attach the EBS volume to all the EC2 instances.

D. Create file systems on Amazon Elastic Block Store (Amazon EBS) volumes that are attached to each EC2 instance. Synchronize the EBS volumes across the different EC2 instances.

Correct Answer: B

Community vote distribution

**Comments**

Josantru **Highly Voted** 1 year, 10 months ago

Correct B.

How is Amazon EFS different than Amazon S3?

Amazon EFS provides shared access to data using a traditional file sharing permissions model and hierarchical directory structure via the NFSv4 protocol. Applications that access data using a standard file system interface provided through the operating system can use Amazon EFS to take advantage of the scalability and reliability of file storage in the cloud without writing any new code or adjusting applications.

Amazon S3 is an object storage platform that uses a simple API for storing and accessing data. Applications that do not require a file system structure and are designed to work with object storage can use Amazon S3 as a massively scalable, durable, low-cost object storage solution.

upvoted 13 times

Gape4 **Most Recent** 11 months, 2 weeks ago

Selected Answer: B

<https://aws.amazon.com/efs/when-to-choose-efs/>

upvoted 2 times

zinabu 1 year, 2 months ago

hierarchical structure =EFS

upvoted 3 times

pentium75 1 year, 5 months ago

Selected Answer: B

Not A because S3 does not allow a "hierarchical directory structure"

Not C because Multi-attach does not work "across two Availability Zones"

Not D because we need "shared", not synchronized, storage.

upvoted 4 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

hierarchical directory structure, read and write rapidly and concurrently to shared storage = Amazon Elastic File System

upvoted 2 times

zits88 10 months ago

Best and most concise explanation here. All true.

upvoted 1 times

potomac 1 year, 7 months ago

Selected Answer: B

Amazon EFS simultaneously supports on-premises servers using a traditional file permissions model, file locking, and hierarchical directory structure through the NFS v4 protocol.

upvoted 2 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

The key reasons:

EFS provides a scalable, high performance NFS file system that can be concurrently accessed from multiple EC2 instances.

It supports the hierarchical directory structure needed by the applications.

EFS is elastic, growing and shrinking automatically as needed.

It can be accessed from instances across AZs, meeting the shared storage requirement.

S3 object storage (option A) lacks the file system semantics needed by the apps.

EBS volumes (options C and D) are attached to a single instance and would require replication and syncing to share across instances.

EFS is purpose-built for this use case of a shared file system across Linux instances and aligns best with the performance, concurrency, and availability needs.

upvoted 4 times

barracouto 1 year, 9 months ago

Selected Answer: B

Going with b

upvoted 2 times

Bennyboy789 1 year, 10 months ago

Selected Answer: B

C and D involve using Amazon EBS volumes, which are block storage. While they can be attached to EC2 instances, they might not provide the same level of shared concurrent access as Amazon EFS. Additionally, synchronizing EBS volumes across different EC2 instances (as in option D) can be complex and error-prone.

Therefore, for a scenario where multiple EC2 instances need to rapidly and concurrently access shared storage with a hierarchical directory structure, Amazon EFS is the best solution.

upvoted 3 times

ukivanlampi 1 year, 10 months ago

Selected Answer: B

s3 is flat structure. EBS multi mount only for same available zone
upvoted 2 times

Dana12345 1 year, 10 months ago

Selected Answer: B

Because Amazon EBS Multi-Attach enables you to attach a single Provisioned IOPS SSD (io1 or io2) volume to multiple instances that are in the same Availability Zone. The infra contains 2 AZ's.
upvoted 2 times

mrsoa 1 year, 10 months ago

Selected Answer: B

B is the correct answer

<https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html>
upvoted 2 times

mrsoa 1 year, 10 months ago

B is the correct answer

<https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html>
upvoted 2 times

RazSteel 1 year, 10 months ago

Selected Answer: C

I think that C is the best option coz io2 can share storage and multi attach.
upvoted 1 times

PLN6302 1 year, 9 months ago

hierarchial directory structure is present in EFS
upvoted 3 times

pentium75 1 year, 5 months ago

Multi-attach does not work "across two Availability Zones".
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #567

Topic 1

A solutions architect is designing a workload that will store hourly energy consumption by business tenants in a building. The sensors will feed a database through HTTP requests that will add up usage for each tenant. The solutions architect must use managed services when possible. The workload will receive more features in the future as the solutions architect adds independent components.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon API Gateway with AWS Lambda functions to receive the data from the sensors, process the data, and store the data in an Amazon DynamoDB table. **Most Voted**
- B. Use an Elastic Load Balancer that is supported by an Auto Scaling group of Amazon EC2 instances to receive and process the data from the sensors. Use an Amazon S3 bucket to store the processed data.
- C. Use Amazon API Gateway with AWS Lambda functions to receive the data from the sensors, process the data, and store the data in a Microsoft SQL Server Express database on an Amazon EC2 instance.
- D. Use an Elastic Load Balancer that is supported by an Auto Scaling group of Amazon EC2 instances to receive and process the data from the sensors. Use an Amazon Elastic File System (Amazon EFS) shared file system to store the processed data.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: A

The key reasons are:

- ° API Gateway removes the need to manage servers to receive the HTTP requests from sensors
- ° Lambda functions provide a serverless compute layer to process data as needed
- ° DynamoDB is a fully managed NoSQL database that scales automatically
- ° This serverless architecture has minimal operational overhead to manage
- ° Options B, C, and D all require managing EC2 instances which increases ops workload
- ° Option C also adds SQL Server admin tasks and licensing costs
- ° Option D uses EFS file storage which requires capacity planning and management

upvoted 6 times

mrsoa **Highly Voted** 1 year, 10 months ago

Selected Answer: A

A seems to be the right answer

upvoted 5 times

AWSSURI **Most Recent** 9 months, 2 weeks ago

Options B,C,D involves unwanted operational overheads due to EC2

So A is the right answer

upvoted 2 times

Mikado211 1 year, 6 months ago

Thinking of that, there is not many questions about IoT Core, but this product could be an excellent answer for the need.

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

Workload runs every hour, must use managed services, more features in the future, LEAST operational overhead = AWS Lambda functions.

HTTP requests, must use managed services, more features in the future, LEAST operational overhead = API Gateway.

Must use managed services, more features in the future, LEAST operational overhead = Amazon DynamoDB.

upvoted 4 times

ersin13 1 year, 10 months ago

key word is "must use managed services when possible" api ,lambda dynamodb are serverless. so answer is A

upvoted 2 times

Kiki_Pass 1 year, 10 months ago

Selected Answer: A

"The workload will receive more features in the future ..." -> DynamoDB

upvoted 4 times

Bmaster 1 year, 10 months ago

A is correct.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #568

Topic 1

A solutions architect is designing the storage architecture for a new web application used for storing and viewing engineering drawings. All application components will be deployed on the AWS infrastructure.

The application design must support caching to minimize the amount of time that users wait for the engineering drawings to load. The application must be able to store petabytes of data.

Which combination of storage and caching should the solutions architect use?

- A. Amazon S3 with Amazon CloudFront **Most Voted**
- B. Amazon S3 Glacier with Amazon ElastiCache
- C. Amazon Elastic Block Store (Amazon EBS) volumes with Amazon CloudFront
- D. AWS Storage Gateway with Amazon ElastiCache

Correct Answer: A

Community vote distribution

A (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 9 months ago

Selected Answer: A

The key reasons are:

S3 provides highly durable and scalable object storage capable of handling petabytes of data cost-effectively. CloudFront can be used to cache S3 content at the edge, minimizing latency for users and speeding up access to the engineering drawings.

The global CloudFront edge network is ideal for caching large amounts of static media like drawings.

EBS provides block storage but lacks the scale and durability of S3 for large media files.

Glacier is cheaper archival storage but has higher latency unsuited for frequent access.

Storage Gateway and ElastiCache may play a role but do not align as well to the main requirements.

upvoted 6 times

Bmaster **Highly Voted** 1 year, 10 months ago

A is correct

upvoted 15 times

upvoted 5 times

Gape4 Most Recent 11 months, 2 weeks ago

Selected Answer: A

I think the answer is A.

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

Petabyte data on AWS infra with high performance

B is Glacier so slow

C EBS for petabyte data doesn't work

D Storage gateway is for on premise connectivity which is not required

upvoted 4 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

Storing and viewing engineering drawings = Amazon S3

Support caching to minimize the amount of time that users wait for the engineering drawings to load = Amazon CloudFront

upvoted 2 times

wsdadasdqwdaw 1 year, 7 months ago

CF caching and S3 supports petabytes data

upvoted 3 times

lemur88 1 year, 9 months ago

Selected Answer: A

CF allows caching

upvoted 2 times

mrsoa 1 year, 10 months ago

Selected Answer: A

The answer seems A:

B : Glacier for archiving

C : i dont think EBS scale to petabytes (I am not sure about that)

D : it incorrect becasueAll application components will be deployed on the AWS infrastrucur

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #569

Topic 1


An Amazon EventBridge rule targets a third-party API. The third-party API has not received any incoming traffic. A solutions architect needs to determine whether the rule conditions are being met and if the rule's target is being invoked.

Which solution will meet these requirements?

- A. Check for metrics in Amazon CloudWatch in the namespace for AWS/Events. **Most Voted**
- B. Review events in the Amazon Simple Queue Service (Amazon SQS) dead-letter queue.
- C. Check for the events in Amazon CloudWatch Logs.
- D. Check the trails in AWS CloudTrail for the EventBridge events.

Correct Answer: A

Community vote distribution



Comments

lemur88 **Highly Voted** 1 year, 3 months ago

Selected Answer: A

<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-monitoring.html>

upvoted 8 times

awsgeek75 **Highly Voted** 11 months ago

Selected Answer: A

"EventBridge sends metrics to Amazon CloudWatch every minute for everything from the number of matched events to the number of times a target is invoked by a rule."

from <https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-monitoring.html>

B: SQS, irrelevant

C: 'Check for events', this wording is confusing but could mean something in wrong context. I would have chosen C if A wasn't an option

D: CloudTrail is for AWS resource monitoring so irrelevant

upvoted 7 times

pentium75 **Most Recent** 11 months, 1 week ago

Selected Answer: A

A per <https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-monitoring.html>

Not B because SQS is not even involved here

Not C because EventBridge sends only metrics, not detailed logs, to CloudWatch

Not D, many fall for CloudTrail supposedly recording "API calls", but this is about calls for the EventBridge API to AWS, not calls to 3rd party APIs by EventBridge.

upvoted 6 times

Min_93 11 months, 2 weeks ago

Selected Answer: C

Option A, "Check for metrics in Amazon CloudWatch in the namespace for AWS/Events," primarily provides aggregated metrics related to EventBridge, but it may not give detailed information about individual events or their specific content. Metrics in CloudWatch can give you an overview of how many events are being processed, but for detailed inspection of events and their conditions, checking CloudWatch Logs (option C) is more appropriate.

CloudWatch Logs allow you to see the actual event data and details, providing a more granular view that is useful for troubleshooting and understanding the specifics of why a third-party API is not receiving incoming traffic.

upvoted 1 times

SHAAHIBHUSHANAWS 1 year ago

A

Events not generating logs in cloudwatch and cloudtrail. only metric data is available.

upvoted 2 times

TariqKipkemei 1 year ago

Selected Answer: D

CloudWatch is a monitoring service for AWS resources and applications. CloudTrail is a web service that records API activity in your AWS account. CloudWatch monitors applications and infrastructure performance in the AWS environment. CloudTrail monitors actions in the AWS environment.

upvoted 1 times

pentium75 11 months, 1 week ago

"API activity", referring to AWS APIs. This would record if someone modifies the EventBridge configuration.

upvoted 2 times

EdenWang 1 year ago

Selected Answer: C

C should be correct, I check in AWS management concole.

upvoted 1 times

potomac 1 year, 1 month ago

Selected Answer: A

shoud be A

upvoted 2 times

ibu007 1 year, 3 months ago

Selected Answer: D

Check the trails in AWS CloudTrail for the EventBridge events.

upvoted 1 times

pentium75 11 months, 1 week ago

I think CloudTrail captures management events (such as modifying the EventBridge configuration)

upvoted 2 times

Eminenza22 1 year, 3 months ago

Selected Answer: C

Amazon CloudWatch Logs is a service that collects and stores logs from Amazon Web Services (AWS) resources. These logs can be used to troubleshoot problems, monitor performance, and audit activity. The other options are incorrect:

Option A: CloudWatch metrics are used to track the performance of AWS resources. They are not used to store events.
Option B: Amazon SQS dead-letter queues are used to store messages that cannot be delivered to their intended recipients. They are not used to store events.
Option D: AWS CloudTrail is a service that records AWS API calls. It can be used to track the activity of EventBridge rules, but it does not store the events themselves.
upvoted 2 times

Eminenza22 1 year, 3 months ago

Errata Corrige
A

EventBridge sends metrics to Amazon CloudWatch every minute for everything from the number of matched events to the number of times a target is invoked by a rule.
<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-monitoring.html>
upvoted 2 times

Eminenza22 1 year, 3 months ago

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/events/CloudWatch-Events-Monitoring-CloudWatch-Metrics.html>
upvoted 2 times

jayce5 1 year, 3 months ago

Selected Answer: D

The answer is D:
"CloudTrail captures API calls made by or on behalf of your AWS account from the EventBridge console and to EventBridge API operations." (<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-logging-monitoring.html>)
upvoted 2 times

pentium75 11 months, 1 week ago

"API calls" to AWS for managing EventBridge. Not "API calls" BY EventBridge to 3rd party APIs.
upvoted 1 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: D

The key reasons:

AWS CloudTrail provides visibility into EventBridge operations by logging API calls made by EventBridge. Checking the CloudTrail trails will show the PutEvents API calls made when EventBridge rules match an event pattern. CloudTrail will also log the Invoke API call when the rule target is triggered. CloudWatch metrics and logs contain runtime performance data but not info on rule evaluation and targeting. SQS dead letter queues collect failed event deliveries but won't provide insights on successful invocations. CloudTrail is purpose-built to log operational events and API activity so it can confirm if the EventBridge rule is being evaluated and triggering the target as expected.
upvoted 2 times

Eminenza22 1 year, 3 months ago

Amazon CloudWatch Logs is a service that collects and stores logs from Amazon Web Services (AWS) resources. These logs can be used to troubleshoot problems, monitor performance, and audit activity.

The other options are incorrect:

Option A: CloudWatch metrics are used to track the performance of AWS resources. They are not used to store events.

Option B: Amazon SQS dead-letter queues are used to store messages that cannot be delivered to their intended recipients.

They are not used to store events.

Option D: AWS CloudTrail is a service that records AWS API calls. It can be used to track the activity of EventBridge rules, but it does not store the events themselves.

upvoted 1 times

Bennyboy789 1 year, 4 months ago

Selected Answer: A

Option A is the most appropriate solution because Amazon EventBridge publishes metrics to Amazon CloudWatch. You can find relevant metrics in the "AWS/Events" namespace, which allows you to monitor the number of events matched by the rule and the number of invocations to the rule's target.

upvoted 6 times

upvoted 0 times

h8er 1 year, 4 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/events/CloudWatch-Events-Monitoring-CloudWatch-Metrics.html>
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #570

Topic 1

A company has a large workload that runs every Friday evening. The workload runs on Amazon EC2 instances that are in two Availability Zones in the us-east-1 Region. Normally, the company must run no more than two instances at all times. However, the company wants to scale up to six instances each Friday to handle a regularly repeating increased workload.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create a reminder in Amazon EventBridge to scale the instances.
- B. Create an Auto Scaling group that has a scheduled action. **Most Voted**
- C. Create an Auto Scaling group that uses manual scaling.
- D. Create an Auto Scaling group that uses automatic scaling.

Correct Answer: B*Community vote distribution*B (100%)**Comments****Bmaster** **Highly Voted** 1 year, 4 months ago

B is correct.

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/ec2-auto-scaling-scheduled-scaling.html>
upvoted 9 times

wizcloudifa **Most Recent** 7 months, 1 week ago**Selected Answer: B**

A - too much operation overhead, manually provisioning the instances after you receive the reminder from eventbridge
B - right answer, as you can scale up the EC2 instances and keep them ready before large overload time
C - too much operation overhead in manually scaling
D - automatic scaling will scale up the instances after some duration after it has encountered the heavy workload traffic, not ideal
upvoted 3 times

TariqKipkemei 1 year ago**Selected Answer: B**

runs every Friday evening = an Auto Scaling group that has a scheduled action
upvoted 3 times

Guru4Cloud 1 year, 3 months ago

Selected Answer: B

The key reasons:

Auto Scaling scheduled actions allow defining specific dates/times to scale out or in. This can be used to scale to 6 instances every Friday evening automatically.

Scheduled scaling removes the need for manual intervention to scale up/down for the workload.

EventBridge reminders and manual scaling require human involvement each week adding overhead.

Automatic scaling responds to demand and may not align perfectly to scale out every Friday without additional tuning.

Scheduled Auto Scaling actions provide the automation needed to scale for the weekly workload without ongoing operational overhead.

upvoted 4 times

Sat897 1 year, 4 months ago

Selected Answer: B

Predicted period.. So schedule the instance

upvoted 4 times

mrsoa 1 year, 4 months ago

Selected Answer: B

B seems to be correct

upvoted 2 times

Deepakin96 1 year, 4 months ago

Selected Answer: B

When we know the run time is Friday, we can schedule the instance to 6

upvoted 3 times

Josantru 1 year, 4 months ago

Correct B.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #571

Topic 1


A company is creating a REST API. The company has strict requirements for the use of TLS. The company requires TLSv1.3 on the API endpoints. The company also requires a specific public third-party certificate authority (CA) to sign the TLS certificate.

Which solution will meet these requirements?

- A. Use a local machine to create a certificate that is signed by the third-party CA. Import the certificate into AWS Certificate Manager (ACM). Create an HTTP API in Amazon API Gateway with a custom domain. Configure the custom domain to use the certificate. **Most Voted**
- B. Create a certificate in AWS Certificate Manager (ACM) that is signed by the third-party CA. Create an HTTP API in Amazon API Gateway with a custom domain. Configure the custom domain to use the certificate.
- C. Use AWS Certificate Manager (ACM) to create a certificate that is signed by the third-party CA. Import the certificate into AWS Certificate Manager (ACM). Create an AWS Lambda function with a Lambda function URL. Configure the Lambda function URL to use the certificate.
- D. Create a certificate in AWS Certificate Manager (ACM) that is signed by the third-party CA. Create an AWS Lambda function with a Lambda function URL. Configure the Lambda function URL to use the certificate.

Correct Answer: A

Community vote distribution



Comments

bjexamprep **Highly Voted** 1 year, 9 months ago

Selected Answer: A

I don't understand why some many people vote B. In ACM, you can either request certificate from Amazon CA or import an existing certificate. There is no option in ACM that allow you to request a certificate that can be signed by third party CA.
upvoted 21 times

markoniz 1 year, 8 months ago

I fully agree
upvoted 6 times

undeadendawd 1 year, 7 months ago

wsuasusquwaw 1 year, 7 months ago

Hmm AWS is saying:

ACM certificates can be used to establish secure communications across the internet or within an internal network. You can request a publicly trusted certificate directly from ACM (an "ACM certificate") or import a publicly trusted certificate issued by a third party. Self-signed certificates are also supported. To provision your organization's internal PKI, you can issue ACM certificates signed by a private certificate authority (CA) created and managed by AWS Private CA. The CA may either reside in your account or be shared with you by a different account.

<https://docs.aws.amazon.com/acm/latest/userguide/gs.html>

upvoted 5 times

pentium75 1 year, 5 months ago

Exactly. You can "import [not create] a publicly trusted certificate issued by a third party".

upvoted 5 times

luiscc Highly Voted 1 year, 10 months ago

Selected Answer: B

AWS Certificate Manager (ACM) is a service that lets you easily provision, manage, and deploy SSL/TLS certificates for use with AWS services and your internal resources. By creating a certificate in ACM that is signed by the third-party CA, the company can meet its requirement for a specific public third-party CA to sign the TLS certificate.

upvoted 9 times

pentium75 1 year, 5 months ago

Sounds like ChatGPT answer, "creating a certificate in ACM that is signed by the third-party CA" is not possible.

upvoted 5 times

emakid Most Recent 11 months, 2 weeks ago

Selected Answer: A

A. Use a local machine to create a certificate that is signed by the third-party CA. Import the certificate into AWS Certificate Manager (ACM). Create an HTTP API in Amazon API Gateway with a custom domain. Configure the custom domain to use the certificate.

Reason:

Custom Certificate: Allows you to use a certificate signed by the third-party CA.

TLSv1.3 Support: API Gateway supports TLSv1.3 for custom domains.

Configuration: You can import the third-party CA certificate into ACM and configure API Gateway to use this certificate with a custom domain.

This approach meets all the specified requirements by allowing the use of a third-party CA-signed certificate and ensuring the API endpoints use TLSv1.3.

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

A is logical answer.

BCD are either misworded here or intentionally confusing. Regardless, you cannot create a cert in ACM that is signed by 3rd party CA. You can only import these certs to ACM.

upvoted 4 times

Shubhi_08 1 year, 5 months ago

Selected Answer: A

We can't create third party certificates in ACM.

upvoted 2 times

foha2012 1 year, 5 months ago

Is this a question from the associate or professional exam ??

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: A

ACM can import, but not create, 3rd party certificates. Leaves only A.

upvoted 2 times

maged123 1 year, 5 months ago

Selected Answer: A

You have already a publicly trusted certificate issued by a third party and you just need to import it in ACM not to create a new one. So, the correct answer is A which is the only one that importing the certificate in ACM while B, C and D are creating a new one.

upvoted 2 times

sparun1607 1 year, 6 months ago

The answer must be A,

You can't create a certificate in ACM, read the below link

<https://docs.aws.amazon.com/acm/latest/userguide/setup.html>

upvoted 2 times

numark 1 year, 6 months ago

Answer is A: Can I import a third-party certificate and use it with AWS services?

Yes. If you want to use a third-party certificate with Amazon CloudFront, Elastic Load Balancing, or Amazon API Gateway, you may import it into ACM using the AWS Management Console, AWS CLI, or ACM APIs. ACM does not manage the renewal process for imported certificates. You can use the AWS Management Console to monitor the expiration dates of an imported certificates and import a new third-party certificate to replace an expiring one.

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

It's 22/Nov/2023 and from the console you can't create a certificate in AWS Certificate Manager (ACM) that is signed by the third-party CA. But you could obtain it externally then import it into ACM.

upvoted 2 times

Tshring 1 year, 6 months ago

Selected Answer: B

Option B meets these requirements:

- API Gateway HTTP APIs support TLS 1.3
- ACM can import certificates signed by third-party CAs
- API Gateway provides REST APIs

upvoted 1 times

pentium75 1 year, 5 months ago

"ACM can import (!) certificates signed by third-party CA", but not create (!) them as B suggests.

upvoted 2 times

NickGordon 1 year, 7 months ago

Selected Answer: A

In ACM you can't create a cert signed by another CA. Dude, try it by yourself. There is no such option!

upvoted 2 times

chen0305_099 1 year, 9 months ago

WHY NOT A?

upvoted 1 times

Guru4Cloud 1 year, 9 months ago

Selected Answer: B

Use ACM to create a certificate signed by the third-party CA. ACM integrates with external CAs.

Create an API Gateway HTTP API with a custom domain name.

Configure the custom domain to use the ACM certificate. API Gateway supports configuring custom domains with ACM certificates.

This allows serving the API over TLS using the required third-party certificate and TLS 1.3 support.
upvoted 2 times

pentium75 1 year, 5 months ago

"ACM integrates with external CAs." no
upvoted 2 times

taustin2 1 year, 9 months ago

Selected Answer: A

You can provide certificates for your integrated AWS services either by issuing them directly with ACM or by importing third-party certificates into the ACM management system.
upvoted 2 times

vini15 1 year, 10 months ago

Should be A.
We need to import third-party certificate to ACM.
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #572

Topic 1

A company runs an application on AWS. The application receives inconsistent amounts of usage. The application uses AWS Direct Connect to connect to an on-premises MySQL-compatible database. The on-premises database consistently uses a minimum of 2 GiB of memory.

The company wants to migrate the on-premises database to a managed AWS service. The company wants to use auto scaling capabilities to manage unexpected workload increases.

Which solution will meet these requirements with the LEAST administrative overhead?

- A. Provision an Amazon DynamoDB database with default read and write capacity settings.
- B. Provision an Amazon Aurora database with a minimum capacity of 1 Aurora capacity unit (ACU).
- C. Provision an Amazon Aurora Serverless v2 database with a minimum capacity of 1 Aurora capacity unit (ACU). **Most Voted**
- D. Provision an Amazon RDS for MySQL database with 2 GiB of memory.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Guru4Cloud** **Highly Voted** 1 year, 9 months ago**Selected Answer: C**

The key reasons:

Aurora Serverless v2 provides auto-scaling so the database can handle inconsistent workloads and spikes automatically without admin intervention.

It can scale down to zero when not in use to minimize costs.

The minimum 1 ACU capacity is sufficient to replace the on-prem 2 GiB database based on the info given.

Serverless capabilities reduce admin overhead for capacity management.

DynamoDB lacks MySQL compatibility and requires more hands-on management.

RDS and provisioned Aurora require manually resizing instances to scale, increasing admin overhead.

upvoted 11 times

dkw2342 1 year, 2 months ago

> It can scale down to zero when not in use to minimize costs.
This part is not correct. Aurora Serverless v1 was able to scale to zero.
upvoted 2 times

kambarami Highly Voted 1 year, 8 months ago

the questions are hard from 500 +
upvoted 7 times

foha2012 1 year, 5 months ago

I dont think these are associate exam questions rather are from AWS professional exam
upvoted 3 times

awsgeek75 1 year, 4 months ago

Yes, I agree. I have been reading the pro questions and these are copy paste. On the bright side, it prepares you for the next step!

upvoted 7 times

emakid Most Recent 11 months, 2 weeks ago

Selected Answer: C

C. Provision an Amazon Aurora Serverless v2 database with a minimum capacity of 1 Aurora capacity unit (ACU).

Suitability: Amazon Aurora Serverless v2 is a good option for applications with variable workloads because it automatically adjusts capacity based on demand. It can handle MySQL-compatible databases and supports auto-scaling. You can set the minimum and maximum capacity based on your needs, making it highly suitable for handling unexpected workload increases with minimal administrative overhead.

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: C

LEAST administrative overhead = Aurora Serverless

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

LEAST administrative overhead = Serverless

upvoted 2 times

ibu007 1 year, 9 months ago

Selected Answer: C

serverless = LEAST overhead

upvoted 3 times

D10SJoker 1 year, 10 months ago

Why not D?

upvoted 2 times

wizcloudifa 1 year, 1 month ago

no autoscaling with RDS

upvoted 2 times

awsgeek75 1 year, 4 months ago

Because "LEAST administrative overhead" is a requirement. RDS configured with mem requirements is an admin overhead

upvoted 2 times

mrsoa 1 year, 10 months ago

Selected Answer: C

C seems to be the right answer

C seems to be the right answer

Instead of provisioning and managing database servers, you specify Aurora capacity units (ACUs). Each ACU is a combination of approximately 2 gigabytes (GB) of memory, corresponding CPU, and networking. Database storage automatically scales from 10 gibibytes (GiB) to 128 tebibytes (TiB), the same as storage in a standard Aurora DB cluster

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-serverless-v1.how-it-works.html>

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-serverless-v2.html>

upvoted 2 times

Bmaster 1 year, 10 months ago

C is correct.

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/aurora-serverless-v2.how-it-works.html#aurora-serverless-v2.how-it-works.capacity>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #573

Topic 1

A company wants to use an event-driven programming model with AWS Lambda. The company wants to reduce startup latency for Lambda functions that run on Java 11. The company does not have strict latency requirements for the applications. The company wants to reduce cold starts and outlier latencies when a function scales up.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure Lambda provisioned concurrency.
- B. Increase the timeout of the Lambda functions.
- C. Increase the memory of the Lambda functions.

D. Configure Lambda SnapStart. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Guru4Cloud **Highly Voted** 1 year, 3 months ago

Selected Answer: D

The key reasons:

SnapStart keeps functions initialized and ready to respond quickly, eliminating cold starts. SnapStart is optimized for applications without aggressive latency needs, reducing costs. It scales automatically to match traffic spikes, eliminating outliers when scaling up. SnapStart is a native Lambda feature with no additional charges, keeping costs low. Provisioned concurrency incurs charges for always-on capacity reserved. More costly than SnapStart. Increasing timeout and memory do not directly improve startup performance like SnapStart.

upvoted 15 times

anikety123 **Highly Voted** 1 year, 3 months ago

Selected Answer: D

Both Lambda SnapStart and provisioned concurrency can reduce cold starts and outlier latencies when a function scales up. SnapStart helps you improve startup performance by up to 10x at no extra cost. Provisioned concurrency keeps functions initialized and ready to respond in double-digit milliseconds. Configuring provisioned concurrency incurs charges to your AWS account. Use provisioned concurrency if your application has strict cold start latency requirements. You can't use both SnapStart

and provisioned concurrency on the same function version.

upvoted 5 times

awsgeek75 Most Recent 11 months ago

Selected Answer: D

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>

"Lambda SnapStart for Java can improve startup performance for latency-sensitive applications by up to 10x at no extra cost, typically with no changes to your function code."

upvoted 5 times

awsgeek75 11 months ago

Also

A: Solves concurrency issues not startup

B is for execution timeout (don't think that possible if I understand the option correctly)

C Memory is not the issue here

upvoted 2 times

TariqKipkemei 1 year ago

Selected Answer: D

Lambda SnapStart it is.

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html#:~:text=RSS-,Lambda%20SnapStart,-for%20Java%20can>

upvoted 2 times

TariqKipkemei 1 year ago

only because its a Java 11 app...if it were any other besides Java I believe Provisioned concurrency could help.

upvoted 2 times

potomac 1 year, 1 month ago

Selected Answer: D

Lambda SnapStart for Java can improve startup performance for latency-sensitive applications by up to 10x at no extra cost, typically with no changes to your function code.

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>

upvoted 3 times

BrijMohan08 1 year, 3 months ago

Selected Answer: D

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>

upvoted 2 times

skyphilip 1 year, 3 months ago

Selected Answer: D

D is correct

Lambda SnapStart for Java can improve startup performance for latency-sensitive applications by up to 10x at no extra cost, typically with no changes to your function code. The largest contributor to startup latency (often referred to as cold start time) is the time that Lambda spends initializing the function, which includes loading the function's code, starting the runtime, and initializing the function code.

With SnapStart, Lambda initializes your function when you publish a function version. Lambda takes a Firecracker microVM snapshot of the memory and disk state of the initialized execution environment, encrypts the snapshot, and caches it for low-latency access. When you invoke the function version for the first time, and as the invocations scale up, Lambda resumes new execution environments from the cached snapshot instead of initializing them from scratch, improving startup latency.

upvoted 2 times

avkya 1 year, 3 months ago

"SnapStart does not support provisioned concurrency, the arm64 architecture, Amazon Elastic File System (Amazon EFS), or ephemeral storage greater than 512 MB." The question says "The company wants to reduce cold starts" This means provisioned concurrency. I'm a little bit confused with D.

upvoted 2 times

upvoted 3 times

Woodlawn5700 1 year, 3 months ago

D

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>

upvoted 2 times

mrsoa 1 year, 4 months ago

Selected Answer: D

D is the answer

Lambda SnapStart for Java can improve startup performance for latency-sensitive applications by up to 10x at no extra cost, typically with no changes to your function code. The largest contributor to startup latency (often referred to as cold start time) is the time that Lambda spends initializing the function, which includes loading the function's code, starting the runtime, and initializing the function code.

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>

upvoted 3 times

Bmaster 1 year, 4 months ago

D is best!!

A is not MOST cost effectly.

lambda snapshot is new feature for lambda.

<https://docs.aws.amazon.com/lambda/latest/dg/snapstart.html>

upvoted 4 times

Bmaster 1 year, 4 months ago

misspell.... lambda snapstart

upvoted 2 times

RaksAWS 1 year, 4 months ago

why not D

It should work

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #574

Topic 1

A financial services company launched a new application that uses an Amazon RDS for MySQL database. The company uses the application to track stock market trends. The company needs to operate the application for only 2 hours at the end of each week. The company needs to optimize the cost of running the database.

Which solution will meet these requirements MOST cost-effectively?

- A. Migrate the existing RDS for MySQL database to an Aurora Serverless v2 MySQL database cluster. **Most Voted**
- B. Migrate the existing RDS for MySQL database to an Aurora MySQL database cluster.
- C. Migrate the existing RDS for MySQL database to an Amazon EC2 instance that runs MySQL. Purchase an instance reservation for the EC2 instance.
- D. Migrate the existing RDS for MySQL database to an Amazon Elastic Container Service (Amazon ECS) cluster that uses MySQL container images to run tasks.

Correct Answer: A*Community vote distribution***Comments****Guru4Cloud** **Highly Voted** 1 year, 9 months ago**Selected Answer: A**

The key reasons are:

Aurora Serverless v2 scales compute capacity automatically based on actual usage, down to zero when not in use. This minimizes costs for intermittent usage.

Since it only runs for 2 hours per week, the application is ideal for a serverless architecture like Aurora Serverless.

Aurora Serverless v2 charges per second when the database is active, unlike RDS which charges hourly.

Aurora Serverless provides higher availability than self-managed MySQL on EC2 or ECS.

Using reserved EC2 instances or ECS still incurs charges when not in use versus the fine-grained scaling of serverless.

Standard Aurora clusters have a minimum capacity unlike the auto-scaling serverless architecture.

upvoted 13 times

dkw2342 1 year, 2 months ago

A is correct, but Aurora Serverless v2 only scales down to 0.5 ACU, not to zero.

upvoted 2 times

pentium75 Highly Voted 1 year, 5 months ago

Selected Answer: A

2 hours per week = Serverless = A. Recommended for "infrequent, intermittent, or unpredictable workloads"

upvoted 5 times

LeonSauveterre Most Recent 6 months ago

Selected Answer: A

A: Aurora Serverless v2 offers the most cost-effective and operationally efficient solution for intermittent workloads like this.

B: over-provisioned and too expensive for this workload

C: Also expensive. You gotta pay for the instance 24/7, regardless of usage.

D: Introduces complexity and operational overhead for database management tasks. Overall, ECS is never easy to handle.

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

B is wrong because Aurora MySQL cluster will just keep on running for the rest of the week and will be costly.

C and D have too much infra bloating so costly

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

Answer is A.

Here are the key distinctions:

Amazon Aurora: provides built-in security, continuous backups, serverless compute, up to 15 read replicas, automated multi-Region replication, and integrations with other AWS services.

Amazon Aurora Serverless: is an on-demand, auto-scaling configuration for Aurora where the database automatically starts up, shuts down, and scales capacity up or down based on your application's needs.

With serverless the db will shut down when not in use.

upvoted 5 times

anikety123 1 year, 9 months ago

Selected Answer: A

Option is A

upvoted 3 times

hachiri 1 year, 9 months ago

Selected Answer: A

Aurora Serverless

- Automated database instantiation and auto-scaling based on actual usage
- Good for infrequent, intermittent or unpredictable workloads
- No capacity planning needed
- Pay per second, can be more cost-effective

upvoted 3 times

vini15 1 year, 10 months ago

will go with A

Amazon Aurora Serverless v2 is suitable for the most demanding, highly variable workloads. For example, your database usage might be heavy for a short period of time, followed by long periods of light activity or no activity at all.

upvoted 3 times

msdnpro 1 year, 10 months ago

Selected Answer: A

"Amazon Aurora Serverless v2 is suitable for the most demanding, highly variable workloads. For example, your database usage might be heavy for a short period of time. followed by long periods of light activity or no activity at all. "