Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #613 *Topic 1*

A company uses Amazon Elastic Kubernetes Service (Amazon EKS) to run a container application. The EKS cluster stores sensitive information in the Kubernetes secrets object. The company wants to ensure that the information is encrypted.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use the container application to encrypt the information by using AWS Key Management Service (AWS KMS).

B. Enable secrets encryption in the EKS cluster by using AWS Key Management Service (AWS KMS). **Most Voted**

C. Implement an AWS Lambda function to encrypt the information by using AWS Key Management Service (AWS KMS).

D. Use AWS Systems Manager Parameter Store to encrypt the information by using AWS Key Management Service (AWS KMS).

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

Guru4Cloud Highly Voted 1 year, 8 months ago

Selected Answer: B

EKS supports encrypting Kubernetes secrets at the cluster level using AWS KMS keys. This provides an automated way to encrypt secrets.
Enabling this feature requires minimal configuration changes to the EKS cluster and no code changes.
Other options like using Lambda functions or modifying the application code to encrypt secrets require additional development effort and overhead.
Systems Manager Parameter Store could store encrypted parameters but does not natively integrate with EKS to encrypt Kubernetes secrets.
The EKS secrets encryption feature leverages AWS KMS without the need to directly call KMS APIs from the application.

upvoted 8 times

KennethNg923 Most Recent 11 months, 3 weeks ago

Selected Answer: B

System manager: irrelevant
Lambda or application: operational overhead
So it will be B secret encryption

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

LEAST operational overhead? = Enable secrets encryption in the EKS cluster

upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: B

https://aws.amazon.com/about-aws/whats-new/2020/03/amazon-eks-adds-envelope-encryption-for-secrets-with-aws-kms/

upvoted 3 times

dilaaziz 1 year, 7 months ago

Selected Answer: B

https://aws.amazon.com/about-aws/whats-new/2020/03/amazon-eks-adds-envelope-encryption-for-secrets-with-aws-kms/

upvoted 2 times

iwannabeawsgod 1 year, 7 months ago

BBBBBBB

upvoted 2 times

taustin2 1 year, 8 months ago

Selected Answer: B

Use KMS. Enable secrets encryption in KMS.

upvoted 3 times

nnecode 1 year, 8 months ago

Selected Answer: B

Enabling secrets encryption in the EKS cluster by using AWS Key Management Service (AWS KMS) is the least operationally overhead way to encrypt the sensitive information in the Kubernetes secrets object.

When you enable secrets encryption in the EKS cluster, AWS KMS encrypts the secrets before they are stored in the EKS cluster. You do not need to make any changes to your container application or implement any additional Lambda functions.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #614 *Topic 1*

A company is designing a new multi-tier web application that consists of the following components:

• Web and application servers that run on Amazon EC2 instances as part of Auto Scaling groups
• An Amazon RDS DB instance for data storage

A solutions architect needs to limit access to the application servers so that only the web servers can access them.

Which solution will meet these requirements?

A. Deploy AWS PrivateLink in front of the application servers. Configure the network ACL to allow only the web servers to access the application servers.

B. Deploy a VPC endpoint in front of the application servers. Configure the security group to allow only the web servers to access the application servers.

C. Deploy a Network Load Balancer with a target group that contains the application servers' Auto Scaling group. Configure the network ACL to allow only the web servers to access the application servers.

D. Deploy an Application Load Balancer with a target group that contains the application servers' Auto Scaling group. Configure the security group to allow only the web servers to access the application servers. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (87%) B (13%)

Comments

Guru4Cloud Highly Voted 1 year, 2 months ago

Selected Answer: D

The key reasons are:

An Application Load Balancer (ALB) allows directing traffic to the application servers and provides access control via security groups.
Security groups act as a firewall at the instance level and can control access to the application servers from the web servers.
Network ACLs work at the subnet level and are less flexible for security groups for instance-level access control.
VPC endpoints are used to provide private access to AWS services, not for access between EC2 instances.

AWS PrivateLink provides private connectivity between VPCs, which is not required in this single VPC scenario.
upvoted 21 times

Dz1990 Most Recent 2 weeks, 3 days ago

Selected Answer: D

Option D is correct because:

ALB is the right tool to put in front of application servers
Security groups are the right tool to control access (not network ACLs)
This creates a proper barrier where only web servers can reach the application servers

Option B is wrong because VPC endpoints are for AWS services, not your own EC2 application servers.
Think of it like this: ALB is like a security guard at a building entrance - only people (web servers) with the right ID can get through to reach the workers (application servers) inside.
upvoted 1 times

Ravan 9 months, 1 week ago

Selected Answer: B

A VPC endpoint is a managed endpoint in your VPC that is connected to a public AWS service. It provides a private connection between your VPC and the service, and it does not require an internet gateway or a NAT device.
The other options do not meet all of the requirements:

Option A: AWS PrivateLink is a service that allows you to connect your VPC to private services that are owned by AWS or by other AWS customers. It is not designed to be used to limit access to resources within the same VPC.
Option C: A Network Load Balancer can be used to distribute traffic across multiple application servers, but it does not provide a way to limit access to the application servers.
Option D: An Application Load Balancer can be used to distribute traffic across multiple application servers, but it does not provide a way to limit access to the application servers.
upvoted 2 times

awsgeek75 10 months, 3 weeks ago

Selected Answer: D

"limit access to the application servers so that only the web servers can access them"
Can be done via NACL or SG
A: Irrelevant as everything is inside the same VPC
B: VPC endpoint are attached to VPC and if you deploy a VPC endpoint like this it will be in front of both app and web server. Language is weird here
C: Potentially a good solution but NACL is allowing on web to app traffic and no response will reach to web servers as NACL have to be configured in both directions
D: ALB in front of ASG will give an internal endpoint which can be secured by SG as recommended. ASG itself is not an endpoint that can be used with SG which is why we need ALB here.

Hence D is correct
upvoted 3 times

TariqKipkemei 1 year ago

Selected Answer: D

Deploy an Application Load Balancer with a target group that contains the application servers' Auto Scaling group. Configure the security group to allow only the web servers to access the application servers
upvoted 4 times

Tekk97 1 year ago

Selected Answer: D

I think B also working. but A company has Auto Scaling groups. D has strategy for Auto Scaling. D is correct
upvoted 2 times

pentium75 11 months, 1 week ago

How do you want to "deploy a VPC endpoint" for a group of EC2 instances that are inside your VPC?
upvoted 2 times

potomac 1 year, 1 month ago

Selected Answer: D

D is correct
upvoted 2 times

iwannabeawsgod 1 year, 1 month ago

Selected Answer: D

Scaling group to Scaling group.
upvoted 3 times

Devsin2000 1 year, 2 months ago

C - ALB is for Web applications only. NLB can be internal / not public
upvoted 1 times

pentium75 11 months, 1 week ago

Both ALB and NLB can be internal or public. NLB works on layer 3 while ALB works on layer 7.

Both ALB and NLB could help here, but C uses a network ACL that's missing the outbound traffic.
upvoted 3 times

taustin2 1 year, 2 months ago

Selected Answer: D

ALB with Security Group is simplest solution.
upvoted 4 times

nnecode 1 year, 2 months ago

Selected Answer: B

A VPC endpoint is a managed endpoint in your VPC that is connected to a public AWS service. It provides a private connection between your VPC and the service, and it does not require an internet gateway or a NAT device.
The other options do not meet all of the requirements:

Option A: AWS PrivateLink is a service that allows you to connect your VPC to private services that are owned by AWS or by other AWS customers. It is not designed to be used to limit access to resources within the same VPC.
Option C: A Network Load Balancer can be used to distribute traffic across multiple application servers, but it does not provide a way to limit access to the application servers.
Option D: An Application Load Balancer can be used to distribute traffic across multiple application servers, but it does not provide a way to limit access to the application servers.
upvoted 4 times

pentium75 11 months, 1 week ago

We don't want to connect "to a public AWS service" but to EC2 instances.

"An Application Load Balancer can be used to distribute traffic across multiple application servers, but it does not provide a way to limit access to the application servers" but the Security Group of the web servers does.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #615 *Topic 1*

A company runs a critical, customer-facing application on Amazon Elastic Kubernetes Service (Amazon EKS). The application has a microservices architecture. The company needs to implement a solution that collects, aggregates, and summarizes metrics and logs from the application in a centralized location.

Which solution meets these requirements?

A. Run the Amazon CloudWatch agent in the existing EKS cluster. View the metrics and logs in the CloudWatch console.

B. Run AWS App Mesh in the existing EKS cluster. View the metrics and logs in the App Mesh console.

C. Configure AWS CloudTrail to capture data events. Query CloudTrail by using Amazon OpenSearch Service.

D. Configure Amazon CloudWatch Container Insights in the existing EKS cluster. View the metrics and logs in the CloudWatch console. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (94%) A (6%)

Comments

pentium75 Highly Voted 11 months, 1 week ago

Selected Answer: D

'Running the Amazon CloudWatch agent in the existing EKS cluster' is called Amazon CloudWatch Container Insights: https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/Container-Insights-setup-metrics.html

upvoted 6 times

potomac Highly Voted 1 year, 1 month ago

Selected Answer: D

Use CloudWatch Container Insights to collect, aggregate, and summarize metrics and logs from your containerized applications and microservices. Container Insights is available for Amazon Elastic Container Service (Amazon ECS), Amazon Elastic Kubernetes Service (Amazon EKS), and Kubernetes platforms on Amazon EC2. Container Insights supports collecting metrics from clusters deployed on AWS Fargate for both Amazon ECS and Amazon EKS.

upvoted 6 times

awsgeek75 Most Recent 11 months ago

Selected Answer: D

https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/ContainerInsights.html
"Use CloudWatch Container Insights to collect, aggregate, and summarize metrics and logs from your containerized applications and microservices."
upvoted 5 times

SHAAHIBHUSHANAWS 1 year ago

Selected Answer: D
https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/ContainerInsights.html
upvoted 4 times

TariqKipkemei 1 year ago

Selected Answer: D

EKS monitoring = Amazon CloudWatch Container Insights
upvoted 3 times

Examanier1217 1 year ago

Selected Answer: A

I have worked on it. A is the right answer
upvoted 1 times

pentium75 11 months, 1 week ago

But 'running the Amazon CloudWatch agent in the existing EKS cluster' is called Container Insights, thus D.

https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/Container-Insights-setup-metrics.html
upvoted 3 times

dilaaziz 1 year, 1 month ago

Selected Answer: D

https://aws.amazon.com/cloudwatch/features/
upvoted 2 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: D

The key reasons are:

CloudWatch Container Insights automatically collects metrics and logs from containers running in EKS clusters. This provides visibility into resource utilization, application performance, and microservice interactions.
The metrics and logs are stored in CloudWatch Logs and CloudWatch metrics for central access.
The CloudWatch console allows querying, filtering, and visualizing the metrics and logs in one centralized place.
upvoted 3 times

ErnShm 1 year, 2 months ago

D

Amazon CloudWatch Application Insights facilitates observability for your applications and underlying AWS resources. It helps you set up the best monitors for your application resources to continuously analyze data for signs of problems with your applications.
upvoted 3 times

taustin2 1 year, 2 months ago

Selected Answer: D

What Cloudwatch Container Insights is for.
upvoted 2 times

kambarami 1 year, 2 months ago

https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/deploy-container-insights-EKS.html
upvoted 2 times

awslearnerin2022 1 year, 2 months ago

Selected Answer: A

Cloudwatch monitors applications and provides metrics. Cloudtrail is used for API activities in the account.

upvoted 1 times

nnecode 1 year, 2 months ago

Selected Answer: D

Amazon CloudWatch Container Insights is a service that collects, aggregates, and summarizes metrics and logs from containerized applications. It is designed to work with Amazon EKS and Kubernetes.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #616 *Topic 1*

A company has deployed its newest product on AWS. The product runs in an Auto Scaling group behind a Network Load Balancer. The company stores the product's objects in an Amazon S3 bucket.

The company recently experienced malicious attacks against its systems. The company needs a solution that continuously monitors for malicious activity in the AWS account, workloads, and access patterns to the S3 bucket. The solution must also report suspicious activity and display the information on a dashboard.

Which solution will meet these requirements?

A. Configure Amazon Macie to monitor and report findings to AWS Config.

B. Configure Amazon Inspector to monitor and report findings to AWS CloudTrail.

C. Configure Amazon GuardDuty to monitor and report findings to AWS Security Hub. **Most Voted**

D. Configure AWS Config to monitor and report findings to Amazon EventBridge.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

Guru4Cloud Highly Voted 1 year, 8 months ago

Selected Answer: C

The key reasons are:

Amazon GuardDuty is a threat detection service that continuously monitors for malicious activity and unauthorized behavior. It analyzes AWS CloudTrail, VPC Flow Logs, and DNS logs.
GuardDuty can detect threats like instance or S3 bucket compromise, malicious IP addresses, or unusual API calls.
Findings can be sent to AWS Security Hub which provides a centralized security dashboard and alerts.
Amazon Macie and Amazon Inspector do not monitor the breadth of activity that GuardDuty does. They focus more on data security and application vulnerabilities respectively.
AWS Config monitors for resource configuration changes, not malicious activity.

upvoted 13 times

MatAlves Most Recent 8 months, 3 weeks ago

Selected Answer: C

- Amazon Inspector = automated vulnerability management service

- Amazon GuardDuty = threat detection service that monitors for malicious activity and anomalous behavior to protect AWS accounts, workloads, and data.

upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: C

" continuously monitors for malicious activity in the AWS account, workloads, and access patterns to the S3 bucket" only guard duty for this purpose in the options

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

Amazon Inspector provides you with security assessments of your applications settings and configurations on your EC2 instances while Amazon GuardDuty helps with analyzing your entire AWS environment for potential threats.
AWS Security Hub is a cloud security posture management service that aggregates alerts, and enables automated remediation.

upvoted 4 times

dilaaziz 1 year, 7 months ago

Selected Answer: C

Guardduty

upvoted 3 times

taustin2 1 year, 8 months ago

Selected Answer: C

What Guard Duty is for.

upvoted 3 times

Guru4Cloud 1 year, 8 months ago

The key reasons are:

Amazon GuardDuty is a threat detection service that continuously monitors for malicious activity and unauthorized behavior. It analyzes AWS CloudTrail, VPC Flow Logs, and DNS logs.
GuardDuty can detect threats like instance or S3 bucket compromise, malicious IP addresses, or unusual API calls.
Findings can be sent to AWS Security Hub which provides a centralized security dashboard and alerts.
Amazon Macie and Amazon Inspector do not monitor the breadth of activity that GuardDuty does. They focus more on data security and application vulnerabilities respectively.
AWS Config monitors for resource configuration changes, not malicious activity.

upvoted 3 times

kambarami 1 year, 8 months ago

Answer is C.

upvoted 2 times

aleariva 1 year, 8 months ago

C is the correct. https://aws.amazon.com/guardduty/

upvoted 2 times

brownie23 1 year, 8 months ago

Answer is C Since Amazon GuardDuty is a threat detection service that continuously monitors for malicious activity and unauthorized behavior to protect your AWS accounts, Amazon Elastic Compute Cloud (EC2) workloads, container applications, Amazon Aurora databases, and data stored in Amazon Simple Storage Service (S3).

upvoted 3 times

awslearnerin2022 1 year, 8 months ago

Selected Answer: C

Gaurd duty is a threat detection service for accounts and workloads.

upvoted 2 times

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #617 *Topic 1*

A company wants to migrate an on-premises data center to AWS. The data center hosts a storage server that stores data in an NFS-based file system. The storage server holds 200 GB of data. The company needs to migrate the data without interruption to existing services. Multiple resources in AWS must be able to access the data by using the NFS protocol.

Which combination of steps will meet these requirements MOST cost-effectively? (Choose two.)

A. Create an Amazon FSx for Lustre file system.

B. Create an Amazon Elastic File System (Amazon EFS) file system. **Most Voted**

C. Create an Amazon S3 bucket to receive the data.

D. Manually use an operating system copy command to push the data into the AWS destination.

E. Install an AWS DataSync agent in the on-premises data center. Use a DataSync task between the on-premises location and AWS. **Most Voted**

Correct Answer: *BE*

Community vote distribution

BE (100%)

Comments

Guru4Cloud Highly Voted 1 year, 2 months ago

Selected Answer: BE

Amazon EFS provides a scalable, high performance NFS file system that can be accessed from multiple resources in AWS.
AWS DataSync can perform the migration from the on-prem NFS server to EFS without interruption to existing services.
This avoids having to manually move the data which could cause downtime. DataSync incrementally syncs changed data.
EFS and DataSync together provide a cost-optimized approach compared to using S3 or FSx, while still meeting the requirements.
Manually copying 200 GB of data to AWS would be slow and risky compared to using DataSync.

upvoted 11 times

awsgeek75 Highly Voted 11 months ago

Selected Answer: BE

A: FSX Lustre is for parallel high performance file storage not NFS
C: S3 is a blob storage, not a file system
D: Too much to copy with a lot of overhead

D: Too much to copy with a lot of overhead
A: NFS maps to EFS and allows NFS protocol for access
E: DataSync solves copy problems without interruptions
upvoted 5 times

dilaaziz Most Recent 1 year, 1 month ago

Selected Answer: BE

https://aws.amazon.com/compare/the-difference-between-nfs-smb/
upvoted 2 times

taustin2 1 year, 2 months ago

Selected Answer: BE

NFS file system = EFS, Use DataSync for the migration with NFS support.
upvoted 5 times

awslearnerin2022 1 year, 2 months ago

Selected Answer: BE

EFS can be accessed by multiple AWS resources.
Datasync allowes NFS migrations.
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #618 *Topic 1*

A company wants to use Amazon FSx for Windows File Server for its Amazon EC2 instances that have an SMB file share mounted as a volume in the us-east-1 Region. The company has a recovery point objective (RPO) of 5 minutes for planned system maintenance or unplanned service disruptions. The company needs to replicate the file system to the us-west-2 Region. The replicated data must not be deleted by any user for 5 years.

Which solution will meet these requirements?

A. Create an FSx for Windows File Server file system in us-east-1 that has a Single-AZ 2 deployment type. Use AWS Backup to create a daily backup plan that includes a backup rule that copies the backup to us-west-2. Configure AWS Backup Vault Lock in compliance mode for a target vault in us-west-2. Configure a minimum duration of 5 years.

B. Create an FSx for Windows File Server file system in us-east-1 that has a Multi-AZ deployment type. Use AWS Backup to create a daily backup plan that includes a backup rule that copies the backup to us-west-2. Configure AWS Backup Vault Lock in governance mode for a target vault in us-west-2. Configure a minimum duration of 5 years.

C. Create an FSx for Windows File Server file system in us-east-1 that has a Multi-AZ deployment type. Use AWS Backup to create a daily backup plan that includes a backup rule that copies the backup to us-west-2. Configure AWS Backup Vault Lock in compliance mode for a target vault in us-west-2. Configure a minimum duration of 5 years. **Most Voted**

D. Create an FSx for Windows File Server file system in us-east-1 that has a Single-AZ 2 deployment type. Use AWS Backup to create a daily backup plan that includes a backup rule that copies the backup to us-west-2. Configure AWS Backup Vault Lock in governance mode for a target vault in us-west-2. Configure a minimum duration of 5 years.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

taustin2 Highly Voted 1 year, 8 months ago

Selected Answer: C

Need to use Compliance Mode, so it's either A or C. RPO leads to Multi-AZ so C.

upvoted 13 times

KennethNg923 11 months, 3 weeks ago

Agree Compliance Mode with Multi-AZ
upvoted 2 times

TariqKipkemei Highly Voted 1 year, 6 months ago

Selected Answer: C

high availability = multi AZ
data must be retained for 5 years = compliance mode
upvoted 5 times

TheLaPlanta 1 year, 2 months ago

No HA was mentioned though. But RPO leads to that, so IDK
upvoted 2 times

dilaaziz Most Recent 1 year, 7 months ago

Selected Answer: C

https://docs.aws.amazon.com/aws-backup/latest/devguide/vault-lock.html
upvoted 3 times

thanhnv142 1 year, 7 months ago

C is correct.
A and C is potential answer because they mention compliance mode. But single AZ is recommended for test and development only. For production workloads, we need multi AZ, which is C
upvoted 3 times

Xin123 1 year, 8 months ago

Selected Answer: C

Trust me bro
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #619 *Topic 1*

A solutions architect is designing a security solution for a company that wants to provide developers with individual AWS accounts through AWS Organizations, while also maintaining standard security controls. Because the individual developers will have AWS account root user-level access to their own accounts, the solutions architect wants to ensure that the mandatory AWS CloudTrail configuration that is applied to new developer accounts is not modified.

Which action meets these requirements?

A. Create an IAM policy that prohibits changes to CloudTrail. and attach it to the root user.

B. Create a new trail in CloudTrail from within the developer accounts with the organization trails option enabled.

C. Create a service control policy (SCP) that prohibits changes to CloudTrail, and attach it the developer accounts.
Most Voted

D. Create a service-linked role for CloudTrail with a policy condition that allows changes only from an Amazon Resource Name (ARN) in the management account.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

taustin2 Highly Voted 1 year, 2 months ago
Selected Answer: C

For Organizations to restrict users in accounts, use an SCP.
upvoted 7 times

Xin123 Highly Voted 1 year, 2 months ago
Selected Answer: C

Organizations + Restricts = SCP
upvoted 7 times

awsgeek75 Most Recent 11 months ago
Selected Answer: C

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps.html
upvoted 4 times

awsgeek75 10 months, 3 weeks ago

C is correct but for my sanity I want to know what D is talking about as it makes no sense to me. Can someone explain?
upvoted 2 times

LeonSauveterre 5 months, 4 weeks ago

The "policy condition" mentioned in option D can never be created. A service-linked role is a type of IAM role that is predefined by AWS and tightly integrated with AWS services, allowing AWS services to perform actions on your behalf. However, service-linked roles are not meant to control or restrict user actions, but simply designed for services to function correctly.
upvoted 1 times

TariqKipkemei 1 year ago

Selected Answer: C

Guardrails = service control policy
upvoted 1 times

Ramdi1 1 year, 2 months ago

Selected Answer: C

C - Use SCP best way
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #620 *Topic 1*

A company is planning to deploy a business-critical application in the AWS Cloud. The application requires durable storage with consistent, low-latency performance.

Which type of storage should a solutions architect recommend to meet these requirements?

A. Instance store volume

B. Amazon ElastiCache for Memcached cluster

C. Provisioned IOPS SSD Amazon Elastic Block Store (Amazon EBS) volume **Most Voted**

D. Throughput Optimized HDD Amazon Elastic Block Store (Amazon EBS) volume

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

taustin2 Highly Voted 1 year, 2 months ago

Selected Answer: C

Durable storage excludes A and B. Low-latency excludes D. Choose C.

upvoted 12 times

awsgeek75 Most Recent 11 months ago

Selected Answer: C

AB are not storage or this purpose
D is HDD so slow by nature
C best fit

upvoted 2 times

TariqKipkemei 1 year ago

Selected Answer: C

durable storage, low-latency performance = Provisioned IOPS SSD Amazon EBS volume

upvoted 3 times

potomac 1 year, 1 month ago

-

Selected Answer: C

Provisioned IOPS SSD — Provides high performance for mission-critical, low-latency, or high-throughput workloads.
Throughput Optimized HDD — A low-cost HDD designed for frequently accessed, throughput-intensive workloads.

upvoted 3 times

dilaaziz 1 year, 1 month ago

Selected Answer: C

https://aws.amazon.com/ebs/volume-types/

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #621 *Topic 1*

An online photo-sharing company stores its photos in an Amazon S3 bucket that exists in the us-west-1 Region. The company needs to store a copy of all new photos in the us-east-1 Region.

Which solution will meet this requirement with the LEAST operational effort?

A. Create a second S3 bucket in us-east-1. Use S3 Cross-Region Replication to copy photos from the existing S3 bucket to the second S3 bucket. **Most Voted**

B. Create a cross-origin resource sharing (CORS) configuration of the existing S3 bucket. Specify us-east-1 in the CORS rule's AllowedOrigin element.

C. Create a second S3 bucket in us-east-1 across multiple Availability Zones. Create an S3 Lifecycle rule to save photos into the second S3 bucket.

D. Create a second S3 bucket in us-east-1. Configure S3 event notifications on object creation and update events to invoke an AWS Lambda function to copy photos from the existing S3 bucket to the second S3 bucket.

Correct Answer: *A*

Community vote distribution

A (86%) D (14%)

Comments

Guru4Cloud Highly Voted 1 year, 8 months ago

Selected Answer: A

S3 Cross-Region Replication handles automatically copying new objects added to the source bucket to the destination bucket in a different region.
It continuously replicates new photos without needing to manually copy files or set up Lambda triggers.
CORS only enables cross-origin access, it does not copy objects.
Using Lifecycle rules or Lambda functions requires custom code and logic to handle the copying.
S3 Cross-Region Replication provides automated replication that minimizes operational overhead.
upvoted 8 times

xBUGx Most Recent 1 year, 3 months ago

Selected Answer: D

All NEW photo, not all photo.
We dont want to copy existing photos

We dont want to copy existing photos
upvoted 3 times

MatAlves 8 months, 3 weeks ago

"There are two types of replication: live replication and on-demand replication.

Live replication – To automatically replicate new and updated objects as they are written to the source bucket, use live replication. Live replication doesn't replicate any objects that existed in the bucket before you set up replication. To replicate objects that existed before you set up replication, use on-demand replication."

https://docs.aws.amazon.com/AmazonS3/latest/userguide/replication.html
upvoted 3 times

TheLaPlanta 1 year, 2 months ago

A does exactly that
upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonS3/latest/userguide/replication.html

To automatically replicate new objects as they are written to the bucket, use live replication, such as Cross-Region Replication (CRR). To replicate existing objects to a different bucket on demand, use S3 Batch Replication.
upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

LEAST operational effort = Cross-Region Replication
upvoted 2 times

dilaaziz 1 year, 7 months ago

Selected Answer: A

https://aws.amazon.com/about-aws/whats-new/2015/03/amazon-s3-introduces-cross-region-replication/
upvoted 3 times

taustin2 1 year, 8 months ago

Selected Answer: A

S3 Cross-Region Replication is least operational overhead.
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #622 *Topic 1*

A company is creating a new web application for its subscribers. The application will consist of a static single page and a persistent database layer. The application will have millions of users for 4 hours in the morning, but the application will have only a few thousand users during the rest of the day. The company's data architects have requested the ability to rapidly evolve their schema.

Which solutions will meet these requirements and provide the MOST scalability? (Choose two.)

A. Deploy Amazon DynamoDB as the database solution. Provision on-demand capacity. **Most Voted**

B. Deploy Amazon Aurora as the database solution. Choose the serverless DB engine mode.

C. Deploy Amazon DynamoDB as the database solution. Ensure that DynamoDB auto scaling is enabled.

D. Deploy the static content into an Amazon S3 bucket. Provision an Amazon CloudFront distribution with the S3 bucket as the origin. **Most Voted**

E. Deploy the web servers for static content across a fleet of Amazon EC2 instances in Auto Scaling groups. Configure the instances to periodically refresh the content from an Amazon Elastic File System (Amazon EFS) volume.

Correct Answer: *AD*

Community vote distribution

AD (57%) | CD (31%) | Other (12%)

Comments

taustin2 Highly Voted 1 year, 8 months ago

Selected Answer: AD

Changing answer to A,D. DynamoDB on-demand is more scalable than DynamoDB auto-scaling.

upvoted 12 times

bogobob Highly Voted 1 year, 6 months ago

Selected Answer: CD

For those answering A over C, the question asks about scalability, but the text says the traffic patterns are known and don't state they will change. Both auto-scaling and on-demand can "scale", but auto-scaling is for known, on-demand is better for unknown traffic patterns. Its likely the "scalability" is more to do with the file hosting (EC2 wouldn't scale well at all vs S3)

upvoted 8 times

pentium75 1 year, 5 months ago

"Most scalability" = A. Might be more expensive, but cost is irrelevant in the question.
upvoted 4 times

bora4motion Most Recent 2 weeks, 4 days ago

Selected Answer: BD

I'm going with BD.
"the company's data architects have requested the ability to rapidly evolve their schema"
upvoted 1 times

JA2018 6 months, 2 weeks ago

Selected Answer: BD

Amazon Aurora Serverless for a highly fluctuating DB workload, and S3 for static webpage
upvoted 1 times

marcosviniciuscb 2 months, 2 weeks ago

Aurora is slower than DynamoDB, Dynamo scales to millions of requests, which is what the company asks for
upvoted 1 times

RamanAgarwal 6 months ago

Cant change schema in Aurora
upvoted 1 times

Ravee_L 10 months, 2 weeks ago

B and D B. Deploy Amazon Aurora as the database solution. Choose the serverless DB engine mode: Aurora Serverless is an on-demand, auto-scaling configuration for Amazon Aurora.
upvoted 3 times

a7md0 11 months, 1 week ago

Selected Answer: AD

on-demand capacity for DynamoDB since traffic is not consistent, and S3 & CloudFront for static files
upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: AD

MOST scalability - DynamoDB On-Demand > Auto Scaling + Static Content host in S3
upvoted 2 times

NSA_Poker 1 year ago

Selected Answer: CD

(A) is incorrect. "On-demand capacity mode is probably best when you have new tables with unknown workloads, unpredictable application traffic." We know the workload upper limit is the total amount of subscribers & we know it's busy in the morning & not so much in the afternoon. Nothing in the question says it bursts in the morning.

(C) is correct. The traffic pattern is known, so prep for the morning & the exact upper boundary is the # of subscribers.
upvoted 2 times

jaswantn 1 year, 3 months ago

For autoscaling we need to know the lower and upper limits. Anh the question says....application will have millions of users for 4 hours in the morning....how many millions , how much upper limit we need to set for to handle this much request?
here we can't have exact estimation for the upper limit in autoscaling. Thus, better option is (A)
upvoted 3 times

jaswantn 1 year, 3 months ago

With autoscaling we can face throttling initially, when there is surge of requests and the load is greater than the scaling upper limit. We can gradually increase the upper limit of autoscaling and would be then able to handle the load in subsequent requests preventing ourself from using OnDemand

requests preventing ourself from using OnDemand.
Thus Option (C) is more scalable as it can handle the both types of load(high & low) in efficient manner.
upvoted 1 times

JA2018 6 months, 2 weeks ago

why not Amazon Aurora Serverless?
upvoted 1 times

1Alpha1 1 year, 4 months ago

Selected Answer: AD

AD vs CD ?
1) Please read the final sentence. Which solutions will meet these requirements and provide the "MOST" scalability?
2) It is not possible to predict an exact boundary based on the number of "millions of users".
So I would choose "AD".
upvoted 4 times

NSA_Poker 1 year ago

The exact boundary is known, the application's subscribers. Leaning towards AC now.
upvoted 1 times

06042022 1 year, 4 months ago

Selected Answer: CD

The traffic pattern is known here.
upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: AD

A: On-demand scaling because the demand changes drastically (millions to thousands)
D: S3 for static page is perfect

B: Aurora is RDMS so not much rapid schema changes (it's subjective and DBA will argue but better options on the table are DynamoDB)
E: Too much work and overhead
upvoted 2 times

awsgeek75 1 year, 4 months ago

To be fair, 4 hours is a strange time duration for burst traffic. 20 hours of low traffic may benefit from auto-scaling's as it is long enough to be called a "depressed" traffic mode in autoscaling config. 4 hours in the morning can also be termed as "sustained period" of burst in autoscaling.

This question is not theoretical, someone who has scaled Dynamo in similar scenarios will be able to answer correctly.
upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: AD

Question asks for "most scalability", not cost optimization. "DynamoDB auto scaling ... modifies provisioned throughput settings only when the actual workload stays elevated or depressed for a sustained period of several minutes. ... This means that provisioned capacity is probably best for you if you have relatively predictable application traffic, run applications whose traffic is consistent, and ramps up or down gradually."
upvoted 5 times

pentium75 1 year, 5 months ago

"The on-demand pricing model is ideal for bursty, new, or unpredictable workloads whose traffic can spike in seconds or minutes, and when under-provisioned capacity would impact the user experience."

Whereas on-demand capacity mode is probably best when you have new tables with unknown workloads, unpredictable application traffic and also if you only want to pay exactly for what you use. The on-demand pricing model is ideal for bursty, new, or unpredictable workloads whose traffic can spike in seconds or minutes, and when under-provisioned capacity would impact the user experience.

https://docs.aws.amazon.com/wellarchitected/latest/serverless-applications-lens/capacity.html

upvoted 2 times

Salilgen 5 months, 1 week ago

This support C not A
upvoted 1 times

Ashhher 1 year, 5 months ago

Selected Answer: BD

I understand the argument between A and C, but why not B?
upvoted 2 times

pentium75 1 year, 5 months ago

"Ability to rapidly evolve their schema" -> NoSQL database, schema changes in transactional databases like RDS are difficult
upvoted 5 times

Derek_G 1 year, 5 months ago

Selected Answer: AD

Provisioned on-demand capacity:
Manual: Requires manual setup and management of capacity.
Cost-Effectiveness: Requires manual estimation of workload, which can result in either excess or insufficient capacity.
Use Case: Suitable for relatively stable workloads with predictable capacity needs.

predictable capacity needs.： 4 hours in the morning,a few thousand users during the rest of the day.
upvoted 2 times

[Removed] 1 year, 5 months ago

Selected Answer: CD

Provisioned mode is more suitable and it is the default.
upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: AD

rapidly evolve their schema, MOST scalability for data layer = DynamoDB with on-demand capacity. on-demand capacity mode automatically enables autoscaling.
MOST scalability for single page app = Amazon CloudFront distribution with the S3 bucket as the origin.
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #623 *Topic 1*

A company uses Amazon API Gateway to manage its REST APIs that third-party service providers access. The company must protect the REST APIs from SQL injection and cross-site scripting attacks.

What is the MOST operationally efficient solution that meets these requirements?

A. Configure AWS Shield.

B. Configure AWS WAF. **Most Voted**

C. Set up API Gateway with an Amazon CloudFront distribution. Configure AWS Shield in CloudFront.

D. Set up API Gateway with an Amazon CloudFront distribution. Configure AWS WAF in CloudFront.

Correct Answer: *B*

Community vote distribution

B (88%) D (12%)

Comments

taustin2 Highly Voted 1 year, 2 months ago

Selected Answer: B

SQL Injection and Cross-Site Scripting = WAF so Either B or D. Both B and D are valid options but the question doesn't indicate a real need for CloudFront, so just use WAF with the API Gateway. Answer is B.

upvoted 14 times

awslearnerin2022 Highly Voted 1 year, 2 months ago

Selected Answer: B

WAF helps with layer 7 attacks like SQL injection and XSS. Shield is helpful for DDOS attacks.

upvoted 8 times

marcosviniciuscb Most Recent 2 months, 1 week ago

Selected Answer: D

CloudFront no se elige por su caché
 No se elige por mejorar la latencia
 Se elige porque es el único sitio donde puedes aplicar AWS WAF si usas REST API (v1)

Entonces:

Aunque CloudFront tenga caché y latencia optimizada, eso no es lo relevante aquí.
Lo importante es que es el único recurso donde puedes colgar el WAF para proteger tu REST API.
upvoted 1 times

marcosviniciuscb 2 months, 2 weeks ago

Option B is technically sound and more straightforward if you only want to protect a specific API.

Option D is more scalable and operationally efficient, especially if you're building a modern architecture with multiple services and want protection from the edge.

The question uses the keyword "MOST operationally efficient," so option D is the best.
upvoted 1 times

Salilgen 5 months, 1 week ago

Selected Answer: B

https://aws.amazon.com/blogs/compute/amazon-api-gateway-adds-support-for-aws-waf/
upvoted 1 times

awsgeek75 11 months ago

Selected Answer: B

WAF is good enough for SQL Injection and Cross Site scripting so A is good
A: AWS Shield (basic) is not for SQL injection
C: Same as A
D: Good solution and will work but it provides extra DDoS protection and caching which is not needed (as we don't know much about the API also)
upvoted 3 times

pentium75 11 months, 1 week ago

Selected Answer: B

Question asks for protection against SQL injection and XSS, both is provided by WAF (option B). D would work too, but it would add another layer (CloudFront) with benefits that nobody asked for (and that would cost money), thus it would IMO be less 'operationally efficient'.
upvoted 3 times

Naijaboy99 11 months, 2 weeks ago

Selected Answer: D

D. Set up API Gateway with an Amazon CloudFront distribution. Configure AWS WAF in CloudFront.

Option A (Configure AWS Shield) is a DDoS protection service but doesn't specifically address SQL injection and cross-site scripting attacks.

Option B (Configure AWS WAF) alone is a valid option, but integrating it with CloudFront (Option D) provides additional benefits like improved performance through caching.

Option C (Set up API Gateway with CloudFront and configure AWS Shield in CloudFront) might provide DDoS protection, but for SQL injection and cross-site scripting, AWS WAF is the more appropriate service.
upvoted 5 times

LeonSauveterre 5 months, 4 weeks ago

In questions like this that mentioned "operationally efficient" or "operationally effective", aim for the indicated jobs only. We don't need additional benefits, so CloudFront is not a necessity here.
upvoted 1 times

TariqKipkemei 1 year ago

Selected Answer: B

SQL injection and cross-site scripting attacks = AWS WAF
upvoted 4 times

potomac 1 year, 1 month ago

Selected Answer: B

B or D
But no need for CloudFront
upvoted 2 times

Sugarbear_01 1 year, 1 month ago

Selected Answer: B

AWS WAF protect agains :
Presence of SQL code that is likely to be malicious (known as SQL injection).
Presence of a script that is likely to be malicious (known as cross-site scripting).

AWS Shield provides protection against distributed denial of service (DDoS) attacks for AWS resources, at the network and transport layers (layer 3 and 4) and the application layer (layer 7).

https://docs.aws.amazon.com/waf/latest/developerguide/what-is-aws-waf.html
upvoted 2 times

thanhnv142 1 year, 1 month ago

Finally, I am here at the end. Thank you guys for your support!
upvoted 5 times

Guru4Cloud 1 year, 2 months ago

Selected Answer: B

B. Configure AWS WAF.
upvoted 5 times

aleariva 1 year, 2 months ago

B is the correct. https://docs.aws.amazon.com/waf/latest/developerguide/classic-web-acl-xss-conditions.html
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #624 *Topic 1*

A company wants to provide users with access to AWS resources. The company has 1,500 users and manages their access to on-premises resources through Active Directory user groups on the corporate network. However, the company does not want users to have to maintain another identity to access the resources. A solutions architect must manage user access to the AWS resources while preserving access to the on-premises resources.

What should the solutions architect do to meet these requirements?

A. Create an IAM user for each user in the company. Attach the appropriate policies to each user.

B. Use Amazon Cognito with an Active Directory user pool. Create roles with the appropriate policies attached.

C. Define cross-account roles with the appropriate policies attached. Map the roles to the Active Directory groups.

D. Configure Security Assertion Markup Language (SAML) 2 0-based federation. Create roles with the appropriate policies attached Map the roles to the Active Directory groups. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (89%) B (11%)

Comments

pentium75 Highly Voted 11 months, 1 week ago

Selected Answer: D

Though you can federate Cognito with Active Directory, Cognito is for providing access to your own applications, NOT to AWS Resources.

upvoted 11 times

tsdsmth Highly Voted 10 months, 4 weeks ago

Selected Answer: D

While Amazon Cognito can integrate with Active Directory, it is more focused on providing identity management for mobile and web applications. In this scenario, where the primary concern is integrating with existing on-premises resources, using SAML-based federation with IAM roles is more appropriate.

upvoted 8 times

Salilgen Most Recent 5 months, 1 week ago

Selected Answer: D

https://aws.amazon.com/blogs/security/aws-federated-authentication-with-active-directory-federation-services-ad-fs/
B is wrong because user pool is in Cognito not in Active Directory
https://aws.amazon.com/blogs/security/how-to-set-up-amazon-cognito-for-federated-authentication-using-azure-ad/
upvoted 1 times

sangavi_vijay 11 months, 2 weeks ago

Selected Answer: B

why its not b?
upvoted 2 times

TariqKipkemei 1 year ago

Selected Answer: D

Use Amazon Cognito via SAML integration. (SAML) is an open federation standard that allows an identity provider (for this case on-prem AD) to authenticate users and pass identity and security information about them to a service provider (for this case AWS).

I will settle for D, because this is definitely required for this to work.
upvoted 5 times

NickGordon 1 year, 1 month ago

Selected Answer: D

D.

An Amazon Cognito user pool is a user directory for WEB and MOBILE app authentication and authorization. So it is not a best option for corporate users.
upvoted 3 times

potomac 1 year, 1 month ago

Selected Answer: D

I think it is D
upvoted 2 times

ahlofan 1 year, 1 month ago

Selected Answer: B

Access to Aws resource -> cognito, then use iam role
SAML or AD -> identity pool
upvoted 2 times

pentium75 11 months, 1 week ago

Cognito is for app users, to authenticate users accessing your apps. Cognito is NOT for granting access to AWS resources.
upvoted 3 times

dilaaziz 1 year, 1 month ago

Selected Answer: D

https://aws.amazon.com/identity/saml/
upvoted 2 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #625 *Topic 1*

A company is hosting a website behind multiple Application Load Balancers. The company has different distribution rights for its content around the world. A solutions architect needs to ensure that users are served the correct content without violating distribution rights.

Which configuration should the solutions architect choose to meet these requirements?

A. Configure Amazon CloudFront with AWS WAF.

B. Configure Application Load Balancers with AWS WAF

C. Configure Amazon Route 53 with a geolocation policy **Most Voted**

D. Configure Amazon Route 53 with a geoproximity routing policy

Correct Answer: *C*

Community vote distribution

C (77%) A (23%)

Comments

potomac Highly Voted 1 year, 7 months ago

Selected Answer: C

Geolocation routing policy — Use when you want to route traffic based on the location of users.

Geo-proximity routing policy — Use when you want to route traffic based on the location of your resources and optionally switch resource traffic at one location to resources elsewhere.

upvoted 14 times

MatAlves 8 months, 3 weeks ago

"Restrict the geographic distribution of your content
You can use geographic restrictions, sometimes known as geo blocking, to prevent users in specific geographic locations from accessing content that you're distributing through an Amazon CloudFront distribution."

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/georestrictions.html

upvoted 2 times

MatAlves 8 months, 3 weeks ago

"Violation to distribution rights" is something you want to handle seriously by actually blocking the access. This can only be accomplished by "OPTION A".
upvoted 1 times

loverduck 9 months, 3 weeks ago

They're not talking about detail like your case
upvoted 1 times

pentium75 1 year, 5 months ago

DNS routing can be easily bypassed, and just routing traffic from different countries to different endpoints does still not restrict what each country can see. It's clearly A.
upvoted 4 times

upliftinghut Highly Voted 1 year, 4 months ago

Selected Answer: C

"You can also use geolocation routing to restrict distribution of content to only the locations in which you have distribution rights"
Link: https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geo.html
upvoted 8 times

EllenLiu Most Recent 5 months, 2 weeks ago

Selected Answer: C

A is incorrect, what is config cloudfront with waf? should be WAF geo match statement only
upvoted 1 times

bujuman 7 months ago

Selected Answer: C

When you use geolocation routing, you can localize your content and present some or all of your website in the language of your users. You can also use geolocation routing to restrict distribution of content to only the locations in which you have distribution rights. Another possible use is for balancing load across endpoints in a predictable, easy-to-manage way, so that each user location is consistently routed to the same endpoint.
upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: A

"Restrict the geographic distribution of your content
You can use geographic restrictions, sometimes known as geo blocking, to prevent users in specific geographic locations from accessing content that you're distributing through an Amazon CloudFront distribution."

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/georestrictions.html
upvoted 1 times

MatAlves 8 months, 3 weeks ago

Actually, it also says:

"You can also use geolocation routing to restrict distribution of content to only the locations in which you have distribution rights."

It can be either C or A, but C seems to address the question wording better.
upvoted 2 times

JA2018 6 months, 2 weeks ago

Agreed
upvoted 1 times

Johnoppong101 9 months, 4 weeks ago

Selected Answer: C

Check line 6 on this documentation.
...You can also use geolocation routing to restrict distribution of content to only the locations in which you have distribution

rights.
https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geo.html
upvoted 2 times

Lin878 11 months, 3 weeks ago

Selected Answer: A

Route 53 can only restrict for geolocation users in this case, it's not for contents. I vote for "A".
upvoted 2 times

FZBianco 1 year ago

Answer is A.
upvoted 1 times

xBUGx 1 year, 2 months ago

Selected Answer: A

I vote for A
upvoted 1 times

Ravan 1 year, 3 months ago

Selected Answer: C

. Configure Amazon Route 53 with a geolocation policy.

By configuring Amazon Route 53 with a geolocation policy, the solutions architect can direct users to different Application Load Balancers based on their geographical location. This allows the company to serve the correct content to users in different regions without violating distribution rights. Geolocation routing policies enable you to route traffic based on the geographic location of your users, ensuring that users are directed to the nearest or most appropriate endpoint based on their location. This solution is suitable for scenarios where content distribution rights vary by region and need to be enforced accordingly.
upvoted 4 times

Pics00094 1 year, 3 months ago

Selected Answer: A

I think it's A
upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

WAF for filtering web traffic based on rules. In this case it may be IP address, geolocation, region. CloudFront for global distribution.

B: Just balances and does not filter
CD: Connects the user to the NEAREST server which is not same as AUTHORISED content
upvoted 1 times

awsgeek75 1 year, 4 months ago

WAF for geo filtering can be configured like this:
https://docs.aws.amazon.com/waf/latest/developerguide/waf-rule-statement-type-geo-match.html

How CloudFront integrates with your WAF rules:
https://docs.aws.amazon.com/waf/latest/developerguide/cloudfront-features.html
upvoted 2 times

Marco_St 1 year, 5 months ago

Selected Answer: A

distributions + restriction of content deleivery target = A
upvoted 1 times

pentium75 1 year, 5 months ago

Selected Answer: A

We want to restrict access by country. People in Spain are allowed to access certain content while people in Portugal are not. A

Route 53 geolocation policy that returns the "nearest" endpoint will not help, because a) the "nearest" endpoint could be identical for multiple countries with different distribution rights and b) it could easily be bypassed.
upvoted 3 times

Salilgen 5 months, 1 week ago

A doesn't talk about geographic restrictions
upvoted 1 times

master9 1 year, 5 months ago

Selected Answer: A

AWS CloudFront supports geographic restrictions, also known as geo-blocking, which can be used to control the distribution of your content based on the geographic location of your viewers.

You can use the CloudFront geographic restrictions feature to either grant permission to your users to access your content only if they're in one of the approved countries on your allowlist, or prevent your users from accessing your content if they're in one of the banned countries on your denylist.

For example, if a request comes from a country where you are not authorized to distribute your content, you can use CloudFront geographic restrictions to block the request.
upvoted 1 times

pentium75 1 year, 5 months ago

Edit: Even though you can specify DNS targets by country, this will not help.
upvoted 2 times

Murtadhaceit 1 year, 5 months ago

Selected Answer: C

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geo.html
upvoted 4 times

ekisako 1 year, 6 months ago

Selected Answer: A

https://repost.aws/knowledge-center/cloudfront-geo-restriction
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #626 *Topic 1*

A company stores its data on premises. The amount of data is growing beyond the company's available capacity.

The company wants to migrate its data from the on-premises location to an Amazon S3 bucket. The company needs a solution that will automatically validate the integrity of the data after the transfer.

Which solution will meet these requirements?

A. Order an AWS Snowball Edge device. Configure the Snowball Edge device to perform the online data transfer to an S3 bucket

B. Deploy an AWS DataSync agent on premises. Configure the DataSync agent to perform the online data transfer to an S3 bucket. **Most Voted**

C. Create an Amazon S3 File Gateway on premises Configure the S3 File Gateway to perform the online data transfer to an S3 bucket

D. Configure an accelerator in Amazon S3 Transfer Acceleration on premises. Configure the accelerator to perform the online data transfer to an S3 bucket.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

TariqKipkemei Highly Voted 1 year, 6 months ago

Selected Answer: B

During a transfer, AWS DataSync always checks the integrity of your data.

https://docs.aws.amazon.com/datasync/latest/userguide/configure-data-verification-options.html

upvoted 12 times

potomac Highly Voted 1 year, 7 months ago

Selected Answer: B

During a transfer, AWS DataSync always checks the integrity of your data, but you can specify how and when this verification happens with the following options: Verify only the data transferred (recommended) – DataSync calculates the checksum of happens with the following options: Verify only the data transferred (recommended) – DataSync calculates the checksum of transferred files and metadata at the source location.

https://docs.aws.amazon.com/datasync/latest/userguide/configure-data-verification-options.html
upvoted 6 times

KennethNg923 Most Recent 11 months, 3 weeks ago

Selected Answer: B

"automatically validate the integrity of the data after the transfer" -> so it should be datasync
upvoted 2 times

dilaaziz 1 year, 7 months ago

Selected Answer: B

https://aws.amazon.com/datasync/faqs/
upvoted 2 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #627 *Topic 1*

A company wants to migrate two DNS servers to AWS. The servers host a total of approximately 200 zones and receive 1 million requests each day on average. The company wants to maximize availability while minimizing the operational overhead that is related to the management of the two servers.

What should a solutions architect recommend to meet these requirements?

A. Create 200 new hosted zones in the Amazon Route 53 console Import zone files. **Most Voted**

B. Launch a single large Amazon EC2 instance Import zone tiles. Configure Amazon CloudWatch alarms and notifications to alert the company about any downtime.

C. Migrate the servers to AWS by using AWS Server Migration Service (AWS SMS). Configure Amazon CloudWatch alarms and notifications to alert the company about any downtime.

D. Launch an Amazon EC2 instance in an Auto Scaling group across two Availability Zones. Import zone files. Set the desired capacity to 1 and the maximum capacity to 3 for the Auto Scaling group. Configure scaling alarms to scale based on CPU utilization.

Correct Answer: *A*

Community vote distribution

A (95%) D (5%)

Comments

awsgeek75 Highly Voted 11 months ago

Selected Answer: A

Key requirement it "maximize availability while minimizing the operational overhead" of 200 zones to process million requests

R53 is designed exactly to do this and supports zone import functionality so literally does the job of their EC2 servers but much better so BCD become "overhead" by default. I doubt D will work.

upvoted 5 times

pentium75 Most Recent 11 months, 1 week ago

Selected Answer: A

B, C and D would not "maximize availability" (not HA) and also not minimize the operational overhead.

upvoted 4 times

TariqKipkemei 1 year ago

Selected Answer: A

'maximize availability while minimizing the operational overhead' = severless = Amazon Route 53

upvoted 4 times

EdenWang 1 year ago

Selected Answer: A

Only A makes sense

upvoted 3 times

NickGordon 1 year, 1 month ago

Selected Answer: A

Should be A

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/migrate-dns-domain-in-use.html

upvoted 4 times

potomac 1 year, 1 month ago

Selected Answer: D

D makes more sense to me

upvoted 1 times

pentium75 11 months, 1 week ago

No, "Desired capacity 1" meaning that usually only 1 server would run, but they want to "maximize availability". And operating EC2 servers would not be "minimizing the operational overhead that is related to the management of the two servers."

upvoted 3 times

awsgeek75 10 months, 3 weeks ago

1 EC2 server for millions of requests?

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #628 *Topic 1*

A global company runs its applications in multiple AWS accounts in AWS Organizations. The company's applications use multipart uploads to upload data to multiple Amazon S3 buckets across AWS Regions. The company wants to report on incomplete multipart uploads for cost compliance purposes.

Which solution will meet these requirements with the LEAST operational overhead?

A. Configure AWS Config with a rule to report the incomplete multipart upload object count.

B. Create a service control policy (SCP) to report the incomplete multipart upload object count.

C. Configure S3 Storage Lens to report the incomplete multipart upload object count. Most Voted

D. Create an S3 Multi-Region Access Point to report the incomplete multipart upload object count.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

warp Highly Voted 1 year, 1 month ago

Selected Answer: C

S3 storage lenses can be used to find incomplete multipart uploads: https://aws.amazon.com/blogs/aws-cloud-financial-management/discovering-and-deleting-incomplete-multipart-uploads-to-lower-amazon-s3-costs/

upvoted 7 times

LocNV Highly Voted 11 months, 2 weeks ago

Selected Answer: C

S3 Storage Lens provides four Cost Efficiency metrics for analyzing incomplete multipart uploads in your S3 buckets. These metrics are free of charge and automatically configured for all S3 Storage Lens dashboards.

Incomplete Multipart Upload Storage Bytes – The total bytes in scope with incomplete multipart uploads
% Incomplete MPU Bytes – The percentage of bytes in scope that are results of incomplete multipart uploads
Incomplete Multipart Upload Object Count – The number of objects in scope that are incomplete multipart uploads
% Incomplete MPU Objects – The percentage of objects in scope that are incomplete multipart uploads
https://aws.amazon.com/blogs/aws-cloud-financial-management/discovering-and-deleting-incomplete-multipart-uploads-to-lower-amazon-s3-costs/

upvoted 5 times

awsgeek75 Most Recent 11 months ago

Selected Answer: C

ABD cannot do any of this so C is the right product for this use case

upvoted 3 times

TariqKipkemei 1 year ago

Selected Answer: C

Amazon S3 Storage Lens is a cloud storage analytics solution with support for AWS Organizations to give you organization-wide visibility into object storage, with point-in-time metrics and trend lines as well as actionable recommendations.

upvoted 4 times

potomac 1 year, 1 month ago

Selected Answer: C

C for sure

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #629 *Topic 1*

A company runs a production database on Amazon RDS for MySQL. The company wants to upgrade the database version for security compliance reasons. Because the database contains critical data, the company wants a quick solution to upgrade and test functionality without losing any data.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an RDS manual snapshot. Upgrade to the new version of Amazon RDS for MySQL.

B. Use native backup and restore. Restore the data to the upgraded new version of Amazon RDS for MySQL.

C. Use AWS Database Migration Service (AWS DMS) to replicate the data to the upgraded new version of Amazon RDS for MySQL.

D. Use Amazon RDS Blue/Green Deployments to deploy and test production changes. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

TariqKipkemei Highly Voted 1 year, 6 months ago

Selected Answer: D

A blue/green deployment copies a production database environment to a separate, synchronized staging environment. You can make changes to the database in the staging environment without affecting the production environment. When you are ready, you can promote the staging environment to be the new production database environment, with downtime typically under one minute.

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/blue-green-deployments.html

upvoted 10 times

warp Highly Voted 1 year, 7 months ago

Selected Answer: D

You can make changes to the RDS DB instances in the green environment without affecting production workloads. For example, you can upgrade the major or minor DB engine version, upgrade the underlying file system configuration, or change database parameters in the staging environment. You can thoroughly test changes in the green environment.
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/blue-green-deployments-overview.html

upvoted 6 times

cjace Most Recent 12 months ago

Option A (Create an RDS manual snapshot and upgrade) is the most straightforward and least operationally intensive method to upgrade your Amazon RDS for MySQL instance while ensuring data safety and allowing thorough testing of application functionality post-upgrade. This approach leverages RDS's snapshot capabilities to provide a reliable rollback mechanism if needed, making it the recommended choice for your scenario.

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: D

Least overhead, only CD qualify and D is actually a managed solution for what is being proposed (hopefully) in C so it's better.

upvoted 4 times

foha2012 1 year, 5 months ago

C works for me

upvoted 1 times

potomac 1 year, 7 months ago

Selected Answer: D

D is the answer

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #630 *Topic 1*

A solutions architect is creating a data processing job that runs once daily and can take up to 2 hours to complete. If the job is interrupted, it has to restart from the beginning.

How should the solutions architect address this issue in the MOST cost-effective manner?

A. Create a script that runs locally on an Amazon EC2 Reserved Instance that is triggered by a cron job.

B. Create an AWS Lambda function triggered by an Amazon EventBridge scheduled event.

C. Use an Amazon Elastic Container Service (Amazon ECS) Fargate task triggered by an Amazon EventBridge scheduled event. Most Voted

D. Use an Amazon Elastic Container Service (Amazon ECS) task running on Amazon EC2 triggered by an Amazon EventBridge scheduled event.

Correct Answer: *C*

Community vote distribution

C (91%) B (9%)

Comments

awsgeek75 Highly Voted 1 year, 4 months ago

Selected Answer: C

A: Nonsense
B: Lambda max running time is 15 mins
D: EC2 is more expensive than Fargate for 2 hours duration as EC2 instance will be billed.

upvoted 10 times

awsgeek75 1 year, 4 months ago

A is also nonsense because an EC2 reserved instance will cost the most for the period when the 2 hour job is not running!

upvoted 4 times

pentium75 Highly Voted 1 year, 5 months ago

Selected Answer: C

Not B because of running time

upvoted 5 times

KennethNg923 Most Recent 11 months, 3 weeks ago

Selected Answer: C

EC2 expensive than Fargate or Lambda, but Lambda has 15 mins limit, so only could choose Fargate for micro service which is C.
I think no one will create script for that by the way

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

AWS Fargate will bill you based on the amount of vCPU, RAM, OS, CPU architecture, and storage that your containerized apps consume while running on EKS or ECS.

upvoted 2 times

cevin93 1 year, 6 months ago

Selected Answer: C

should be C

upvoted 3 times

Alex1atd 1 year, 6 months ago

Selected Answer: C

Lambda function have a limit timeout about 15 minutes, so cannot be B.
Answer is C

upvoted 3 times

hungta 1 year, 6 months ago

Selected Answer: C

Lamda function have a limit timeout about 15 minutes

upvoted 2 times

cciesam 1 year, 7 months ago

Selected Answer: B

I think it should be B. Considering the Cost.

upvoted 3 times

Murtadhaceit 1 year, 5 months ago

Lambda times out after 15 minutes. This job requires a 2-hour time without interruption block. So, definitely not B.

upvoted 5 times

zhdetn 1 year, 6 months ago

Lambda Maximum execution time: 900 seconds (15 minutes)

upvoted 6 times

potomac 1 year, 7 months ago

Selected Answer: C

I guess it is C

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #631 *Topic 1*

A social media company wants to store its database of user profiles, relationships, and interactions in the AWS Cloud. The company needs an application to monitor any changes in the database. The application needs to analyze the relationships between the data entities and to provide recommendations to users.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use Amazon Neptune to store the information. Use Amazon Kinesis Data Streams to process changes in the database.

B. Use Amazon Neptune to store the information. Use Neptune Streams to process changes in the database. **Most Voted**

C. Use Amazon Quantum Ledger Database (Amazon QLDB) to store the information. Use Amazon Kinesis Data Streams to process changes in the database.

D. Use Amazon Quantum Ledger Database (Amazon QLDB) to store the information. Use Neptune Streams to process changes in the database.

Correct Answer: *B*

Community vote distribution

B (88%) C (12%)

Comments

pentium75 Highly Voted 1 year, 5 months ago

Selected Answer: B

Relationships between entities = Graph data = Neptune

upvoted 7 times

awsgeek75 1 year, 4 months ago

Also, Neptune Streams can monitor changes in the data and create a changelog

upvoted 6 times

MatAlves 8 months, 3 weeks ago

I come looking for one of you guys comments. Nice.

upvoted 2 times

TariqKinkemei Highly Voted 1 year, 6 months ago

Selected Answer: B

Amazon Neptune Database is a serverless graph database designed for superior scalability and availability. Neptune Database provides built-in security, continuous backups, and integrations with other AWS services. Suitable for social media. With the Neptune Streams feature, you can generate a complete sequence of change-log entries that record every change made to your graph data as it happens.

upvoted 7 times

FlyingHawk Most Recent 4 months, 1 week ago

Selected Answer: B

Neptune is optimized for storing billions of relationships and querying the graph with milliseconds latence. Neptune Streams logs every change to your graph as it happens, in the order that it is made, in a fully managed way.
Amazon Quantum Ledger Database (QLDB) is a purpose-built ledger database that provides a complete and cryptographically verifiable history of all changes made to your application data.

upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: B

Normally Amazon Quantum Ledger Database use in blockchain DB more. So i will go for B using Neptune and Neptune Stream for relationship between entities.

upvoted 2 times

haci 1 year, 3 months ago

Selected Answer: C

Amazon QLDB tracks and maintains a sequential history of every application data change using an immutable and transparent log. It trusts the integrity of your data. Built-in cryptographic authentication provides third-party verification of data changes. QLDB ACID transactions can create accurate, event-driven systems with support for real-time streaming to Amazon Kinesis.

upvoted 1 times

NickGordon 1 year, 7 months ago

Selected Answer: B

B

Social network -> Graph Structure -> Neptune

upvoted 3 times

ekisako 1 year, 7 months ago

Selected Answer: B

Keyword: analyze the relationships
With Amazon Neptune, you can create sophisticated, interactive graph applications that can query billions of relationships in milliseconds.

https://aws.amazon.com/neptune/features/

upvoted 5 times

potomac 1 year, 7 months ago

Selected Answer: C

Amazon Neptune is primarily used for managing highly connected graph data, making it well-suited for graph-based applications.

In contrast, Amazon QLDB is designed for applications that require an immutable and auditable transaction history to ensure data integrity.

upvoted 3 times

pentium75 1 year, 5 months ago

Exactly, thus B. "Relationships between the data entities" is "graph data".

upvoted 2 times

warp 1 year, 7 months ago

Selected Answer: B

Neptune is a graph type database and Neptune streams provides view on changes into the database:
https://docs.aws.amazon.com/neptune/latest/userguide/streams.html

upvoted 3 times

AF_1221 1 year, 7 months ago

C is the correct answer
provides a well-suited, managed, and scalable solution for storing and monitoring the database with the least operational overhead, meeting the requirements of the social media company.

upvoted 2 times

awsgeek75 1 year, 4 months ago

AQLB is like a blockchain database. Are you sure this is the correct option for graph data?

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #632 *Topic 1*

A company is creating a new application that will store a large amount of data. The data will be analyzed hourly and will be modified by several Amazon EC2 Linux instances that are deployed across multiple Availability Zones. The needed amount of storage space will continue to grow for the next 6 months.

Which storage solution should a solutions architect recommend to meet these requirements?

A. Store the data in Amazon S3 Glacier. Update the S3 Glacier vault policy to allow access to the application instances.

B. Store the data in an Amazon Elastic Block Store (Amazon EBS) volume. Mount the EBS volume on the application instances.

C. Store the data in an Amazon Elastic File System (Amazon EFS) file system. Mount the file system on the application instances. **Most Voted**

D. Store the data in an Amazon Elastic Block Store (Amazon EBS) Provisioned IOPS volume shared between the application instances.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

TariqKipkemei Highly Voted 6 months, 1 week ago

Selected Answer: C

Multiple Linux instances = Amazon Elastic File System (Amazon EFS) with multiple mount targets.

upvoted 6 times

AF_1221 Highly Voted 7 months, 1 week ago

C is correct
Shared File System: Amazon EFS allows multiple Amazon EC2 instances to mount the same file system simultaneously, making it easy for multiple instances to access and modify the data concurrently.

upvoted 5 times

Salilgen Most Recent 5 months, 1 week ago

Selected Answer: C

S3 Glacier (flexible retrieval) is not compatible with hourly job
EBS can be shared only in same AZ (B and D are out)
upvoted 2 times

potomac 7 months ago

Selected Answer: C

C is correct
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #633 *Topic 1*

A company manages an application that stores data on an Amazon RDS for PostgreSQL Multi-AZ DB instance. Increases in traffic are causing performance problems. The company determines that database queries are the primary reason for the slow performance.

What should a solutions architect do to improve the application's performance?

A. Serve read traffic from the Multi-AZ standby replica.

B. Configure the DB instance to use Transfer Acceleration.

C. Create a read replica from the source DB instance. Serve read traffic from the read replica. Most Voted

D. Use Amazon Kinesis Data Firehose between the application and Amazon RDS to increase the concurrency of database requests.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

TariqKipkemei Highly Voted 1 year ago

Selected Answer: C

A Multi-AZ DB instance Creates a primary DB instance with one standby DB instance in a different Availability Zone. Using a Multi-AZ DB instance provides high availability, but the standby DB instance doesn't support connections for read workloads. Therefore you will need to create a read replica from the source DB instance then serve read traffic from the read replica.

upvoted 7 times

awsgeek75 Highly Voted 11 months ago

Selected Answer: C

Read replica split for read traffic will distribute the overall load and improve the performance.

A: Standby replica cannot serve traffic (Correct me if I am wrong here)
B: Transfer Accelerator is to speed up S3 traffic. Not the case here
D: Kiensis will increase concurrency but won't solve the DB performance issues

upvoted 5 times

potomac Most Recent 1 year, 1 month ago

potomac 1 year, 1 month ago

Selected Answer: C

you can't read from the standby DB instance. If applications require more read capacity, you should create or add additional read replicas.

upvoted 3 times

warp 1 year, 1 month ago

Selected Answer: C

After you create a read replica from a source DB instance, the source becomes the primary DB instance. When you make updates to the primary DB instance, Amazon RDS copies them asynchronously to the read replica.
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_ReadRepl.html

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #634 *Topic 1*

A company collects 10 GB of telemetry data daily from various machines. The company stores the data in an Amazon S3 bucket in a source data account.

The company has hired several consulting agencies to use this data for analysis. Each agency needs read access to the data for its analysts. The company must share the data from the source data account by choosing a solution that maximizes security and operational efficiency.

Which solution will meet these requirements?

A. Configure S3 global tables to replicate data for each agency.

B. Make the S3 bucket public for a limited time. Inform only the agencies.

C. Configure cross-account access for the S3 bucket to the accounts that the agencies own. **Most Voted**

D. Set up an IAM user for each analyst in the source data account. Grant each user access to the S3 bucket.

Correct Answer: *C*

Community vote distribution

C (96%) D

Comments

pentium75 Highly Voted 11 months, 1 week ago

Selected Answer: C

A doesn't exist
B is a big "hell no"
D is a bad practice, even with IAM you'd use groups

upvoted 9 times

xBUGx Highly Voted 9 months ago

What if other agencies don't have an aws account?

upvoted 7 times

chickenmf 8 months, 4 weeks ago

then we politely tell them "no."

then we politely tell them "no."
upvoted 9 times

awsgeek75 Most Recent 11 months ago

Selected Answer: C

Others have given reason by ABD are wrong. In case you need it, here is an AWS example exercise of understanding option C

https://docs.aws.amazon.com/AmazonS3/latest/userguide/example-walkthroughs-managing-access-example2.html
upvoted 5 times

TariqKipkemei 1 year ago

Selected Answer: C

With cross-account bucket permissions Account A—can grant another AWS account, Account B, permission to access its resources such as buckets and objects. Account B can then delegate those permissions to users in its account.

https://docs.aws.amazon.com/AmazonS3/latest/userguide/example-walkthroughs-managing-access-example2.html#:~:text=4%3A%20Clean%20up-,An%20AWS%20account,-%E2%80%94for%20example%2C%20Account
upvoted 3 times

NickGordon 1 year, 1 month ago

Selected Answer: C

C is the best answer
upvoted 3 times

cciesam 1 year, 1 month ago

Selected Answer: D

C may not correct as it's doesn't say if the analyst are using AWS services
upvoted 1 times

awsgeek75 11 months ago

"consulting agencies" means some companies which may have one or more analysts each. Making IAM users for each individual to manage permissions is not well-architected. You would at least create groups and assign it to users.
D will work as it is possible but it won't minimize "operational efficiency"
upvoted 2 times

NickGordon 1 year, 1 month ago

in that case, an analyst user group should be created and the access should be assigned to the group. So C is better
upvoted 3 times

potomac 1 year, 1 month ago

Selected Answer: C

I think it is C
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #635 *Topic 1*

A company uses Amazon FSx for NetApp ONTAP in its primary AWS Region for CIFS and NFS file shares. Applications that run on Amazon EC2 instances access the file shares. The company needs a storage disaster recovery (DR) solution in a secondary Region. The data that is replicated in the secondary Region needs to be accessed by using the same protocols as the primary Region.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an AWS Lambda function to copy the data to an Amazon S3 bucket. Replicate the S3 bucket to the secondary Region.

B. Create a backup of the FSx for ONTAP volumes by using AWS Backup. Copy the volumes to the secondary Region. Create a new FSx for ONTAP instance from the backup.

C. Create an FSx for ONTAP instance in the secondary Region. Use NetApp SnapMirror to replicate data from the primary Region to the secondary Region. **Most Voted**

D. Create an Amazon Elastic File System (Amazon EFS) volume. Migrate the current data to the volume. Replicate the volume to the secondary Region.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

awsgeek75 Highly Voted 1 year, 4 months ago

Selected Answer: C

This is a very rare usage scenario so here are the docs related to the product:
https://docs.aws.amazon.com/fsx/latest/ONTAPGuide/scheduled-replication.html

AD: Not compatible solutions
B: Either wrongly worded or missing something but if I read it correctly, it means just take a backup and restore whereas the question is about continuous replication. If B was scheduled then it would have made sense
C is correct as SnapMirror is a managed solution to replicate data

upvoted 9 times

KennethNg923 11 months, 3 weeks ago

Agree, thus it said "replicated in the secondary Region" only, not create a backup infrastructure in other region for failover, i don't think we need to use Backup feature
upvoted 2 times

pentium75 Most Recent 1 year, 5 months ago

Selected Answer: C

Not A, no access with CIFS (SMB) or NFS
Not B, one-time copy
Not D, EFS does not offer SMB
upvoted 3 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

C
https://aws.amazon.com/blogs/storage/cross-region-disaster-recovery-with-amazon-fsx-for-netapp-ontap/
upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

Amazon FSx for NetApp ONTAP supports NetApp SnapMirror, a replication technology that you can use to replicate data between two ONTAP file systems. You can configure automatic NetApp SnapMirror replication of your data to another Amazon FSx for NetApp ONTAP file system, including a file system in another AWS Region. If needed, you can fail over your applications and users to use the other Amazon FSx for NetApp ONTAP file system. With SnapMirror, you can configure replication with a Recovery Point Objective (RPO) of as low as 5 minutes, and a Recovery Time Objective (RTO) in single-digit minutes. You can configure SnapMirror using the ONTAP CLI or REST API.
upvoted 3 times

Oblako 1 year, 6 months ago

Selected Answer: C

SnapMirror enables you to configure replication with an RPO of as low as five minutes, and an RTO in single digit minutes. It is the recommended solution for DR when using FSx for ONTAP: https://aws.amazon.com/blogs/storage/cross-region-disaster-recovery-with-amazon-fsx-for-netapp-ontap/
upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: C

You can use NetApp SnapMirror to schedule periodic replication of your FSx for ONTAP file system to or from a second file system. This capability is available for both in-Region and cross-Region deployments.

https://docs.aws.amazon.com/fsx/latest/ONTAPGuide/scheduled-replication.html
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #636 *Topic 1*

A development team is creating an event-based application that uses AWS Lambda functions. Events will be generated when files are added to an Amazon S3 bucket. The development team currently has Amazon Simple Notification Service (Amazon SNS) configured as the event target from Amazon S3.

What should a solutions architect do to process the events from Amazon S3 in a scalable way?

A. Create an SNS subscription that processes the event in Amazon Elastic Container Service (Amazon ECS) before the event runs in Lambda.

B. Create an SNS subscription that processes the event in Amazon Elastic Kubernetes Service (Amazon EKS) before the event runs in Lambda

C. Create an SNS subscription that sends the event to Amazon Simple Queue Service (Amazon SQS). Configure the SOS queue to trigger a Lambda function. **Most Voted**

D. Create an SNS subscription that sends the event to AWS Server Migration Service (AWS SMS). Configure the Lambda function to poll from the SMS event.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

potomac Highly Voted 1 year, 7 months ago

Selected Answer: C

Amazon SQS is designed for event-driven and scalable message processing. It can handle large volumes of messages and automatically scales based on the incoming workload. This allows for better load distribution and scaling as compared to direct Lambda invocation.

upvoted 6 times

TariqKipkemei Highly Voted 1 year, 6 months ago

Selected Answer: C

scalable service = serverless = Amazon SQS implemented with FAN-OUT.
However SQS is a pull based event distribution service, it does not trigger other services.
C is the closest option.

upvoted 5 times

KennethNg923 Most Recent 11 months, 3 weeks ago

Selected Answer: C

SQS + SNS standard solution

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: C

AB are way too complicated to scale without more specifics (no idea about number of events)
D SMS is not for this, it's for server migrations
C SNS is notified on file creation in S3. SNS publishes to SQS which can scale according to the input load automatically.
Lambda execution can scale a lot when attached to SQS.

ABC have scaling limits each but C's scaling limit is much better than AB

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #637 *Topic 1*

A solutions architect is designing a new service behind Amazon API Gateway. The request patterns for the service will be unpredictable and can change suddenly from 0 requests to over 500 per second. The total size of the data that needs to be persisted in a backend database is currently less than 1 GB with unpredictable future growth. Data can be queried using simple key-value requests.

Which combination ofAWS services would meet these requirements? (Choose two.)

A. AWS Fargate
B. AWS Lambda **Most Voted**
C. Amazon DynamoDB **Most Voted**
D. Amazon EC2 Auto Scaling
E. MySQL-compatible Amazon Aurora

Correct Answer: *BC*

Community vote distribution

BC (100%)

Comments

potomac Highly Voted 1 year, 7 months ago
Selected Answer: BC

B and C
upvoted 11 times

TariqKipkemei Highly Voted 1 year, 6 months ago
Selected Answer: BC

Scalable, unpredictable request patterns = AWS Lambda
Scalable, key-value data = Amazon DynamoDB
upvoted 11 times

KennethNg923 Most Recent 11 months, 3 weeks ago
Selected Answer: BC

Auto Scaling cannot handle "suddenly from 0 requests to over 500 per second", use Lambda and Dynamo which for Key-value pair.

upvoted 2 times

Phi143 1 year, 2 months ago

Why not AC? Size of the data has unpredictable future growth and Lambda may not be able to handle it.

upvoted 2 times

JA2018 6 months, 2 weeks ago

Lambda is fully managed by AWS, fully scalable to meet actual spikes in demand

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: BC

Unpredictable scaling of API load = Lambda + SPI Gateway
Unpredictable growth of key/value DB = DynamoDB

Fargate behind API requires EKS/ECS setup which is not suitable for 0-500 varying load. Same with EC2 autoscaling.
Aurora MySQL is not ideal for key/value and is better suited for relational databases

upvoted 4 times

Ashhher 1 year, 5 months ago

Selected Answer: BC

why not Fargate?

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #638 *Topic 1*

A company collects and shares research data with the company's employees all over the world. The company wants to collect and store the data in an Amazon S3 bucket and process the data in the AWS Cloud. The company will share the data with the company's employees. The company needs a secure solution in the AWS Cloud that minimizes operational overhead.

Which solution will meet these requirements?

A. Use an AWS Lambda function to create an S3 presigned URL. Instruct employees to use the URL. **Most Voted**

B. Create an IAM user for each employee. Create an IAM policy for each employee to allow S3 access. Instruct employees to use the AWS Management Console.

C. Create an S3 File Gateway. Create a share for uploading and a share for downloading. Allow employees to mount shares on their local computers to use S3 File Gateway.

D. Configure AWS Transfer Family SFTP endpoints. Select the custom identity provider options. Use AWS Secrets Manager to manage the user credentials Instruct employees to use Transfer Family.

Correct Answer: *A*

Community vote distribution

A (38%) C (38%) D (24%)

Comments

t0nx Highly Voted 1 year, 6 months ago

Selected Answer: D

AWS Transfer Family (Option D)

By configuring AWS Transfer Family SFTP endpoints, you can provide a secure and convenient way for employees to access and transfer data to and from the S3 bucket.
Using custom identity provider options allows you to integrate with existing identity systems, and AWS Secrets Manager can be used to manage user credentials securely.

A suggests using an AWS Lambda function to create an S3 presigned URL. While this can work, it involves manual generation of URLs and sharing them, which may not be as scalable or user-friendly.

B suggests creating an IAM user for each employee with IAM policies for S3 access. This involves more operational overhead, as managing IAM users for each employee can be cumbersome and less scalable.

C suggests using an S3 File Gateway. While this can work, it introduces additional components and may not be as straightforward or as efficient as using AWS Transfer Family for SFTP access.
upvoted 16 times

pentium75 1 year, 5 months ago

"Use AWS Secrets Manager to manage the user credentials", so manage separate credentials for every user in Secrets Manager? And "instruct employees to use Transfer Family", actually Transfer Family is the server component, employees would use an SFTP client.
upvoted 8 times

xxichlas 11 months, 2 weeks ago

https://docs.aws.amazon.com/secretsmanager/latest/userguide/integrating_FTPlong.html
upvoted 2 times

pentium75 Highly Voted 1 year, 5 months ago

Selected Answer: C

Not A - S3 presigned URLs are temporary (max. 7 days); you'd need to create a new URL at least every 7 days and "instruct employees" to use it. Definitely NOT 'minimizing operational overhead'.

Not B - "Instruct employees to use the AWS Management Console", using Management console to up- and download files is complex

Not D - Secrets Manager is not for managing user credentials, and employees would not "use Transfer Family", they would use an (S)FTP client to access the files.

C grants simple access for up/downloading, no operational overhead.
upvoted 13 times

bora4motion 1 month, 1 week ago

but then you will have to manage the S3 File Gateway in the long run, which could be a VM, patching etc. I am tempted to go with D.
upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Agree, Use an AWS Lambda function to create an S3 presigned URL for 7 days limits, create URL every 7 days have operational overhead more than use Secret Manager
upvoted 5 times

awsgeek75 1 year, 4 months ago

Glad that someone else also sees what I see in this question!
upvoted 4 times

ODRAMIREZ Most Recent 1 week, 1 day ago

Selected Answer: C

I think is C
upvoted 1 times

Dz1990 2 weeks, 3 days ago

Selected Answer: A

Option A is correct because:

Lowest operational overhead - just one Lambda function vs managing many users/systems
Secure - presigned URLs expire automatically
Simple for employees - just click a link, no special software needed
Works globally - URLs work from anywhere
Scalable - handles any number of employees easily

Think of it like this: Instead of giving everyone a key to your house (IAM users), you give them a temporary visitor pass (presigned URL) that expires after a set time.

Option D: Transfer Family SFTP

How it works: SFTP service with custom identity provider
Security: □ Secure
Operational overhead: □ MEDIUM-HIGH - managing SFTP endpoints, user credentials in Secrets Manager
Global access: □ Works globally
upvoted 1 times

srcntpc 3 weeks, 2 days ago

Selected Answer: A

Why A is Correct:
S3 presigned URLs allow secure, time-limited access to S3 objects without needing to manage user credentials or IAM permissions for each employee.
AWS Lambda function can dynamically generate presigned URLs based on request parameters or business logic.
Simple, secure, and minimizes operational overhead.
upvoted 1 times

zdi561 4 months ago

Selected Answer: A

Sharing S3 access with one who has no amazon credential is to provide presigned URL. C is not right becasue S3 File Gateway is for on-premise access(corporate) to S3.
upvoted 1 times

FlyingHawk 5 months, 2 weeks ago

Selected Answer: C

I will vote for C. For A, the S3 presigned URL is expired in 7 days, you need to implement a process to allow user send an access request, then validate the user, then trigger the lamdba to generate the resigned URL and send it back to the user. For D, using AWS secret manager to manage the user credentials does not sound practical unless it uses Active Directory.
upvoted 1 times

bora4motion 1 month, 1 week ago

With C, in the long run you will have to manage the s3 gateway which is more for on prem use. So this means you have to patch it, update it, reboot it etc. I am tempted to go with D. With pre-signed url you will have to lean on users to regenerate the links, D - yes, AD for auth and that's it.
upvoted 1 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: A

S3 File Gateway is designed for hybrid cloud use cases where on-premises applications need to interact with S3. It introduces additional infrastructure complexity and costs, making it less suitable for the described requirements.
upvoted 1 times

LeonSauveterre 5 months, 3 weeks ago

Selected Answer: C

A - Definitely not ideal. Presigned URLs have a maximum expiration time of 7 days, which might be limiting for ongoing or long-term data sharing. If employees require frequent or dynamic access to multiple files (which is likely the case because this is a research-type company), they'd need a new URL each time.

B - This is cumbersome and tiresome.

C - This is better, but still not ideal for global employees unless the gateway is highly distributed.

D - Secrets Manager stores sensitive information like passwords or API keys but is not a user directory, so they cannot "manage" credentials.
upvoted 2 times

Rhydian25 11 months, 2 weeks ago

Selected Answer: C

It is not operationally eficient to manage, for example, 1000 signed URLs or user credentials. In addition, it is sometimes dificult to instruct that many people.

It's easier to create an S3 File Gateway and allow the users to mount it locally to access the bucket.

It could be D if the answer said to use IAM roles instead of managing user credentials in Secrets Manager
upvoted 5 times

MandAsh 11 months, 3 weeks ago

Selected Answer: C

but they didnt mention access in for daily use of occasional.
If its occasional A works well but its permanant thing them mouting drive is solution.
upvoted 4 times

stalk98 1 year ago

Selected Answer: D

i think is d
upvoted 1 times

TwinSpark 1 year ago

Selected Answer: C

Less operational overheade is C
https://docs.aws.amazon.com/filegateway/latest/files3/GettingStartedAccessFileShare.html
on client pc is easily mounted. I remain with some doubts but i will go for C
upvoted 3 times

alawada 1 year, 2 months ago

i would go with A
upvoted 1 times

seetpt 1 year, 3 months ago

Selected Answer: D

D seems right
upvoted 1 times

Ravan 1 year, 3 months ago

Selected Answer: A

A. Use an AWS Lambda function to create an S3 presigned URL.

This solution meets the requirements by providing a secure way for employees to access the data stored in the Amazon S3 bucket. Here's how it works:

When an employee needs to access the data, they request access from the company's system.
The company's system triggers an AWS Lambda function.
The Lambda function generates a presigned URL with a limited validity period.
The employee uses the presigned URL to access the data directly from the S3 bucket.
Once the presigned URL expires, access to the data is no longer possible, enhancing security.
This solution minimizes operational overhead because it leverages AWS Lambda, which is a fully managed service. There is no need to manage servers or infrastructure, and the solution provides a secure and temporary access mechanism for sharing data stored in Amazon S3.
upvoted 9 times

FlyingHawk 5 months, 2 weeks ago

if you use this method, when a user request to access, you need to implement the request validation, aka the employee validation, such as integrated with Active Directory.
upvoted 1 times

NayeraB 1 year, 3 months ago

I legitimately get worried every time we have a tie
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #639

Topic 1

A company is building a new furniture inventory application. The company has deployed the application on a fleet ofAmazon EC2 instances across multiple Availability Zones. The EC2 instances run behind an Application Load Balancer (ALB) in their VPC.

A solutions architect has observed that incoming traffic seems to favor one EC2 instance, resulting in latency for some requests.

What should the solutions architect do to resolve this issue?

A. Disable session affinity (sticky sessions) on the ALB **Most Voted**

B. Replace the ALB with a Network Load Balancer

C. Increase the number of EC2 instances in each Availability Zone

D. Adjust the frequency of the health checks on the ALB's target group

Correct Answer: *A*

Community vote distribution

A (87%) B (13%)

Comments

KennethNg923 11 months, 3 weeks ago

Selected Answer: A

"favor one EC2 instance" it is because you enable the sticky session feature, so you have to disable it

upvoted 2 times

1Alpha1 1 year, 4 months ago

Selected Answer: A

Answer: *A*
Enabling stickiness may bring imbalance to the load over the backend EC2 instances since sticky sessions help the same client to always redirect to the same instance behind a load balancer.

upvoted 3 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

The question is too vague. Doesn't say much about the application or EC2 instance setup. So:
If you assume that application uses session management then A is correct.
If you think application is crashing then D is correct for health checks
If you don't assume anything about the application then B is also correct
SMH, I'll go with B... happy to discuss
upvoted 2 times

awsgeek75 1 year, 4 months ago

I'm not entirely happy with any choice but since others have chosen A, I'm just throwing B for discussion.
upvoted 2 times

MikeSWA 1 year, 5 months ago

what about c?
it actually helps distribute traffic equally across instances in all enabled AZs.
upvoted 2 times

mr123dd 1 year, 4 months ago

nope, if the sticky season is on, no matter how many instances you have in AZ or region, it will only send traffic to you favor session
upvoted 2 times

evelynsun 1 year, 5 months ago

it's A!!
Session affinity is a feature of the Application Load Balancer that keeps client requests on the same EC2 instance for the duration of the session. This can cause latency issues if one EC2 instance is overloaded while others are not, as the overloaded instance will handle all subsequent requests until it is taken offline.

To resolve this issue, the solutions architect should disable session affinity on the ALB. This can be done by setting the "Session affinity" parameter to "Off" in the ALB's configuration.

Disabling session affinity will cause the ALB to distribute requests across all EC2 instances in the target group, rather than keeping them on a single instance. This will help to balance the load and reduce latency for all requests.
upvoted 3 times

awsgeek75 1 year, 4 months ago

I agree with A but it assumes that ALB has session affinity enabled and app doesn't require it. What if the EC2 instances are running an application that requires session affinity? I think the question is missing some important context
upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

Disable session affinity (sticky sessions) on the ALB
upvoted 2 times

NickGordon 1 year, 7 months ago

Selected Answer: A

A

https://repost.aws/knowledge-center/elb-fix-unequal-traffic-routing
upvoted 3 times

awsgeek75 1 year, 4 months ago

The same article says to check health of instances. This makes D as a good candidate too.
"Available healthy instances aren't evenly distributed across Availability Zones."
upvoted 4 times

potomac 1 year, 7 months ago

Selected Answer: A

A makes more sense than others
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #640 *Topic 1*

A company has an application workflow that uses an AWS Lambda function to download and decrypt files from Amazon S3. These files are encrypted using AWS Key Management Service (AWS KMS) keys. A solutions architect needs to design a solution that will ensure the required permissions are set correctly.

Which combination of actions accomplish this? (Choose two.)

A. Attach the kms:decrypt permission to the Lambda function's resource policy

B. Grant the decrypt permission for the Lambda IAM role in the KMS key's policy **Most Voted**

C. Grant the decrypt permission for the Lambda resource policy in the KMS key's policy.

D. Create a new IAM policy with the kms:decrypt permission and attach the policy to the Lambda function.

E. Create a new IAM role with the kms:decrypt permission and attach the execution role to the Lambda function. **Most Voted**

Correct Answer: *BE*

Community vote distribution

BE (83%) Other (17%)

Comments

NickGordon Highly Voted 1 year, 7 months ago

Selected Answer: BE

BE is right.

The key policy has to be modified to give lambda execution role access. You can't set another resource policy as principle. So C is not right

upvoted 9 times

vaibhavianup09 Most Recent 4 weeks ago

Selected Answer: BD

B. Grant the decrypt permission for the Lambda IAM role in the KMS key's policy

KMS key policies control who can use the key. The Lambda role must be explicitly allowed in the key policy for decryption to work.

D. Create a new IAM policy with the kms:decrypt permission and attach the policy to the Lambda function

The Lambda function (more precisely, its execution role) must have the kms:Decrypt permission in its IAM policy to invoke the KMS service.
upvoted 1 times

MaxMingxing 5 months, 1 week ago

BE, IAM role act as the agent between lambda function and KMS
upvoted 1 times

1166ae3 11 months, 1 week ago
Selected Answer: BD

E is wrong, AWS Lambda function can hold only one IAM role. This role is known as the execution role. What we should do is: creating an IAM policy that allows the kms:Decrypt action and attach it to the Lambda function's execution role.
upvoted 2 times

cjace 12 months ago

B D - The combination of Option B (Grant the decrypt permission for the Lambda IAM role in the KMS key's policy) and Option D (Create a new IAM policy with the kms
permission and attach the policy to the Lambda function) ensures that both the IAM role used by the Lambda function and the KMS key policy are correctly configured to allow decryption of the files. This setup meets the security requirements and ensures the Lambda function can perform its tasks without issues.
upvoted 1 times

wizcloudifa 1 year ago
Selected Answer: BE

when it comes to permissions look for the "IAM ROLE" word, lambda would need a role to decrypt the s3 object, only roles can be attached to a function not policies
upvoted 3 times

1Alpha1 1 year, 4 months ago
Selected Answer: BE

B. Grant the decrypt permission for the Lambda ***IAM ROLE*** in the KMS key's policy
E. Create a new ***IAM ROLE*** with the kms:decrypt permission and attach the execution role to the Lambda function.
upvoted 4 times

awsgeek75 1 year, 4 months ago
Selected Answer: BE

AC are resource policy, i.e. who can use lambda.
https://docs.aws.amazon.com/lambda/latest/dg/access-control-resource-based.html
D: The wording is confusing so it sort of sounds as if it is correct but you cannot attach a policy to a function.
upvoted 2 times

pentium75 1 year, 5 months ago
Selected Answer: BE

Not A and C because they are about function's "resource policy" which controls who can manage the function, NOT what the function can do.
Not D because you attach an IAM policy to an IAM principal, not to a Lambda function.
upvoted 4 times

TariqKipkemei 1 year, 6 months ago
Selected Answer: BE

Create a new IAM role with the kms:decrypt permission and attach the execution role to the Lambda function then grant the decrypt permission for the Lambda IAM role in the KMS key's policy
upvoted 3 times

louisaok 1 year, 7 months ago
Selected Answer: CE

CE is right
upvoted 1 times

pentium75 1 year, 5 months ago

No, the "Lambda resource policy" is about who can manage the Lambda function

upvoted 2 times

potomac 1 year, 7 months ago

Selected Answer: DE

DE?
Create an IAM role for the Lambda function that also grants decryption permission to the S3 bucket.
Configure the IAM role as the Lambda functions execution role.

To use an IAM policy to control access to a KMS key, the key policy for the KMS key must give the account permission to use IAM policies.

https://repost.aws/knowledge-center/lambda-execution-role-s3-bucket
https://docs.aws.amazon.com/kms/latest/developerguide/iam-policies.html

upvoted 1 times

potomac 1 year, 7 months ago

change to CE
C. Grant the decrypt permission for the Lambda resource policy in the KMS key's policy.
E. Create a new IAM role with the kms:decrypt permission and attach the execution role to the Lambda function.

https://docs.aws.amazon.com/lambda/latest/dg/access-control-resource-based.html
https://docs.aws.amazon.com/kms/latest/developerguide/key-policies.html

upvoted 2 times

pentium75 1 year, 5 months ago

C is about the "Lambda resource policy", who can manage the function.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #641 *Topic 1*

A company wants to monitor its AWS costs for financial review. The cloud operations team is designing an architecture in the AWS Organizations management account to query AWS Cost and Usage Reports for all member accounts. The team must run this query once a month and provide a detailed analysis of the bill.

Which solution is the MOST scalable and cost-effective way to meet these requirements?

A. Enable Cost and Usage Reports in the management account. Deliver reports to Amazon Kinesis. Use Amazon EMR for analysis.

B. Enable Cost and Usage Reports in the management account. Deliver the reports to Amazon S3 Use Amazon Athena for analysis. **Most Voted**

C. Enable Cost and Usage Reports for member accounts. Deliver the reports to Amazon S3 Use Amazon Redshift for analysis.

D. Enable Cost and Usage Reports for member accounts. Deliver the reports to Amazon Kinesis. Use Amazon QuickSight tor analysis.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

NickGordon Highly Voted 7 months ago

Selected Answer: B

B

https://aws.amazon.com/blogs/big-data/analyze-amazon-s3-storage-costs-using-aws-cost-and-usage-reports-amazon-s3-inventory-and-amazon-athena/

upvoted 6 times

LeonSauveterre Most Recent 5 months, 3 weeks ago

Selected Answer: B

A - EMR is powerful for large-scale data processing, but is overkill for a monthly query since neither Kinesis nor EMR is cost-effective for infrequent processing.

B - Cost and Usage Reports are already optimized for S3 delivery, and Athena directly integrates with S3.
C & D - Delivering reports for each member account creates redundant data. That mean extra time, extra work to do manually, extra money to cost, and extra account to manage.

upvoted 1 times

TariqKipkemei 6 months ago

Selected Answer: B

Scalable and cost-effective way = Enable Cost and Usage Reports in the management account. Deliver the reports to Amazon S3 Use Amazon Athena for analysis

upvoted 4 times

potomac 7 months ago

Selected Answer: B

B
once a month

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #642 *Topic 1*

A company wants to run a gaming application on Amazon EC2 instances that are part of an Auto Scaling group in the AWS Cloud. The application will transmit data by using UDP packets. The company wants to ensure that the application can scale out and in as traffic increases and decreases.

What should a solutions architect do to meet these requirements?

A. Attach a Network Load Balancer to the Auto Scaling group. **Most Voted**

B. Attach an Application Load Balancer to the Auto Scaling group.

C. Deploy an Amazon Route 53 record set with a weighted policy to route traffic appropriately.

D. Deploy a NAT instance that is configured with port forwarding to the EC2 instances in the Auto Scaling group.

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

Sugarbear_01 Highly Voted 1 year, 7 months ago

Selected Answer: A

https://docs.aws.amazon.com/autoscaling/ec2/userguide/autoscaling-load-balancer.html

upvoted 11 times

KennethNg923 Most Recent 11 months, 3 weeks ago

Selected Answer: A

UDP packets can scale out and in

upvoted 2 times

awsgeek75 1 year, 4 months ago

Selected Answer: A

UDP can only be monitored by NLB.
ALB is for application layer (HTTP etc)
R53 is DNS
NAT is for port forwarding/address translation etc which is not going to help with scaling

A is correct

upvoted 4 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

UDP packets = Network Load Balancer

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #643 *Topic 1*

A company runs several websites on AWS for its different brands. Each website generates tens of gigabytes of web traffic logs each day. A solutions architect needs to design a scalable solution to give the company's developers the ability to analyze traffic patterns across all the company's websites. This analysis by the developers will occur on demand once a week over the course of several months. The solution must support queries with standard SQL.

Which solution will meet these requirements MOST cost-effectively?

A. Store the logs in Amazon S3. Use Amazon Athena tor analysis. **Most Voted**

B. Store the logs in Amazon RDS. Use a database client for analysis.

C. Store the logs in Amazon OpenSearch Service. Use OpenSearch Service for analysis.

D. Store the logs in an Amazon EMR cluster Use a supported open-source framework for SQL-based analysis.

Correct Answer: *A*

Community vote distribution

A (95%) C (5%)

Comments

awsgeek75 Highly Voted 1 year, 4 months ago

Selected Answer: A

Scalable + "The solution must support queries with standard SQL" = A
B not scalable
C OpenSearch is like ElasticSearch so does not support SQL syntax
D EMR is processing not storage. Map-Reduce can use SQL like syntax but this option does not solve scalable storage issues. You normally run EMR on some stored data

upvoted 6 times

cedser8 1 year, 3 months ago

OpenSearch can support SQL queries, https://docs.aws.amazon.com/opensearch-service/latest/developerguide/sql-support.html

upvoted 2 times

tch Most Recent 4 months, 2 weeks ago

Selected Answer: A

Amazon Athena is an interactive analytics service that makes it simple to analyze data in Amazon Simple Storage Service (S3) using SQL. Athena is serverless, so there is no infrastructure to set up or manage, and you can start analyzing data immediately. You don't even need to load your data into Athena; it works directly with data stored in Amazon S3. Amazon Athena for SQL uses Trino and Presto with full standard SQL support and works with various standard data formats, including CSV, JSON, Apache ORC, Apache Parquet, and Apache Avro. Athena for Apache Spark supports SQL and allows you to use Apache Spark, an open-source, distributed processing system used for big data workloads. To get started, log in to the Athena Management Console and start interacting with your data using the query editor or notebooks.

upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: A

standard SQL + analyze traffic

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: A

Difficult question because both A and C meet the requirements. (OpenSearch does "support queries with standard SQL".)

Still, native S3 storage is slightly cheaper than storage for OpenSearch. Also, Athena does not incur additional cost while OpenSearch does. Question asks for cost efficiency, thus A.

D is out, not only because of the cost but also because you do not 'store logs in (!) an Amazon EMR cluster'; you can use (!) an EMR cluster to analyze data that is stored elsewhere.

upvoted 3 times

pentium75 1 year, 5 months ago

And descriptions of both products, Athena as well as OpenSearch, state that you can use them to "analyze" data.

upvoted 3 times

[Removed] 1 year, 5 months ago

Selected Answer: C

https://docs.aws.amazon.com/opensearch-service/latest/developerguide/cold-storage.html

upvoted 1 times

pentium75 1 year, 5 months ago

Seems that even cold storage is still more expensive than S3.

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: A

solution must support queries with standard SQL = Amazon S3 with Athena

upvoted 4 times

NickGordon 1 year, 7 months ago

Selected Answer: A

A, most cost effective

upvoted 3 times

potomac 1 year, 7 months ago

Selected Answer: A

option D (using Amazon EMR with an open-source framework) may be overkill for the relatively simple SQL-based analysis.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #644 *Topic 1*

An international company has a subdomain for each country that the company operates in. The subdomains are formatted as example.com, country1.example.com, and country2.example.com. The company's workloads are behind an Application Load Balancer. The company wants to encrypt the website data that is in transit.

Which combination of steps will meet these requirements? (Choose two.)

A. Use the AWS Certificate Manager (ACM) console to request a public certificate for the apex top domain example com and a wildcard certificate for *.example.com. **Most Voted**

B. Use the AWS Certificate Manager (ACM) console to request a private certificate for the apex top domain example.com and a wildcard certificate for *.example.com.

C. Use the AWS Certificate Manager (ACM) console to request a public and private certificate for the apex top domain example.com.

D. Validate domain ownership by email address. Switch to DNS validation by adding the required DNS records to the DNS provider.

E. Validate domain ownership for the domain by adding the required DNS records to the DNS provider. **Most Voted**

Correct Answer: *AE*

Community vote distribution

AE (100%)

Comments

awsgeek75 Highly Voted 11 months ago

Selected Answer: AE

B is private certificate so won't help as that is for internal use
C is for apex domain only and won't help with wildcard domain
A is correct

DE are both doable as per these articles

D: https://docs.aws.amazon.com/acm/latest/userguide/dns-validation.html
E: https://docs.aws.amazon.com/acm/latest/userguide/domain-ownership-validation.html

D is less applicable because it does not say if R53 is being used for DNS. You only validate ownership to R53

C makes more sense as it applies to both R53 and other DNS providers

upvoted 9 times

TariqKipkemei Highly Voted 1 year ago

Selected Answer: AE

Validate domain ownership for the domain by adding the required DNS records to the DNS provider then use the AWS Certificate Manager (ACM) console to request a public certificate for the apex top domain example com and a wildcard certificate for *.example.com

upvoted 5 times

cciesam Most Recent 1 year, 1 month ago

Selected Answer: AE

AE correct

upvoted 2 times

potomac 1 year, 1 month ago

Selected Answer: AE

BCD are wrong

upvoted 3 times

t0nx 1 year ago

Why E and not D ?

upvoted 1 times

Cyberkayu 11 months, 3 weeks ago

need to put A-record and CNAME in public DNS record to proof you are the legal owner of the domain name.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #645 *Topic 1*

A company is required to use cryptographic keys in its on-premises key manager. The key manager is outside of the AWS Cloud because of regulatory and compliance requirements. The company wants to manage encryption and decryption by using cryptographic keys that are retained outside of the AWS Cloud and that support a variety of external key managers from different vendors.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use AWS CloudHSM key store backed by a CloudHSM cluster.

B. Use an AWS Key Management Service (AWS KMS) external key store backed by an external key manager. Most Voted

C. Use the default AWS Key Management Service (AWS KMS) managed key store.

D. Use a custom key store backed by an AWS CloudHSM cluster.

Correct Answer: *B*

Community vote distribution

B (95%) A (5%)

Comments

pentium75 Highly Voted 11 months, 1 week ago

Selected Answer: B

Keys are supposed to be managed "outside of the AWS cloud", thus A, C and D are out.

upvoted 7 times

potomac Highly Voted 1 year, 1 month ago

Selected Answer: B

https://docs.aws.amazon.com/kms/latest/developerguide/keystore-external.html

upvoted 5 times

evelynsun Most Recent 11 months, 3 weeks ago

Selected Answer: A

it's A.
This solution is the LEAST operational overhead because it does not require the company to manage any infrastructure or software outside of the AWS Cloud. The AWS CloudHSM key store is managed by AWS, and the company can use it to store

and manage its cryptographic keys without having to worry about the underlying infrastructure or software. The CloudHSM cluster is managed by AWS, and the company can use it to create and manage its cryptographic keys without having to worry about the hardware or software.

the AWS CloudHSM key store can also be used for external key managers. The AWS CloudHSM key store is a managed key store that is backed by an AWS CloudHSM cluster. The AWS CloudHSM cluster is a managed service that is provided by AWS.
upvoted 1 times

pentium75 11 months, 1 week ago

"The AWS CloudHSM key store is managed by AWS" which is exactly what this company does NOT want.
upvoted 6 times

evelynsun 11 months, 3 weeks ago

it's A.
This solution is the LEAST operational overhead because it does not require the company to manage any infrastructure or software outside of the AWS Cloud. The AWS CloudHSM key store is managed by AWS, and the company can use it to store and manage its cryptographic keys without having to worry about the underlying infrastructure or software. The CloudHSM cluster is managed by AWS, and the company can use it to create and manage its cryptographic keys without having to worry about the hardware or software.

the AWS CloudHSM key store can also be used for external key managers. The AWS CloudHSM key store is a managed key store that is backed by an AWS CloudHSM cluster. The AWS CloudHSM cluster is a managed service that is provided by AWS.
upvoted 1 times

pentium75 11 months, 1 week ago

"The AWS CloudHSM key store is managed by AWS" which is exactly what this company does NOT want.
upvoted 3 times

SHAAHIBHUSHANAWS 1 year ago

B
https://docs.aws.amazon.com/kms/latest/developerguide/keystore-external.html
upvoted 4 times

TariqKipkemei 1 year ago

Selected Answer: B

https://docs.aws.amazon.com/kms/latest/developerguide/keystore-external.html#:~:text=Document%20history-,External%20key%20stores,-PDF
upvoted 3 times

1rob 1 year ago

Selected Answer: B

Answer A does not comply because aws cloudHSM is within aws
Answer B is the correct answer because the company is required to use its on-premises key manager. Following https://docs.aws.amazon.com/kms/latest/developerguide/custom-key-store-overview.html gives :An external key store is an AWS KMS custom key store backed by an external key manager outside of AWS that you own and control.(...)

Answer C and D are both solutions in the aws cloud so that does not fit.
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #646 *Topic 1*

A solutions architect needs to host a high performance computing (HPC) workload in the AWS Cloud. The workload will run on hundreds of Amazon EC2 instances and will require parallel access to a shared file system to enable distributed processing of large datasets. Datasets will be accessed across multiple instances simultaneously. The workload requires access latency within 1 ms. After processing has completed, engineers will need access to the dataset for manual postprocessing.

Which solution will meet these requirements?

A. Use Amazon Elastic File System (Amazon EFS) as a shared file system. Access the dataset from Amazon EFS.

B. Mount an Amazon S3 bucket to serve as the shared file system. Perform postprocessing directly from the S3 bucket.

C. Use Amazon FSx for Lustre as a shared file system. Link the file system to an Amazon S3 bucket for postprocessing.
Most Voted

D. Configure AWS Resource Access Manager to share an Amazon S3 bucket so that it can be mounted to all instances for processing and postprocessing.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

potomac Highly Voted 1 year, 7 months ago

Selected Answer: C

Amazon FSx for Lustre is a fully managed, high-performance file system optimized for HPC workloads. It is designed to deliver sub-millisecond latencies and high throughput, making it ideal for applications that require parallel access to shared storage, such as simulations and data analytics.

upvoted 8 times

KennethNg923 Most Recent 11 months, 3 weeks ago

Selected Answer: C

host a high performance computing (HPC) workload -> FSx lustre

upvoted 2 times

zinabu 1 year, 2 months ago

FSx luster for HPC

upvoted 2 times

pentium75 1 year, 5 months ago

Selected Answer: C

EFS could meet the latency requirement for most (!) read (!) operations, but this is not enough here. FSx for Lustre ist specifically designed for HPC.

upvoted 3 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: C

high performance computing (HPC) workloads, shared file system= Amazon FSx for Lustre

upvoted 4 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #647 *Topic 1*

A gaming company is building an application with Voice over IP capabilities. The application will serve traffic to users across the world. The application needs to be highly available with an automated failover across AWS Regions. The company wants to minimize the latency of users without relying on IP address caching on user devices.

What should a solutions architect do to meet these requirements?

A. Use AWS Global Accelerator with health checks. **Most Voted**

B. Use Amazon Route 53 with a geolocation routing policy.

C. Create an Amazon CloudFront distribution that includes multiple origins.

D. Create an Application Load Balancer that uses path-based routing.

Correct Answer: *A*

Community vote distribution

A (97%) C

Comments

potomac Highly Voted 1 year, 7 months ago

Selected Answer: A

Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover.

upvoted 10 times

pentium75 Highly Voted 1 year, 5 months ago

Selected Answer: A

A - does exactly what is required
Not B - Would rely on DNS caching (as it should not)
Not C - CloudFront is not for VoIP
Not D - ALB does not address any of the issues and would not support VoIP

upvoted 7 times

KennethNg923 Most Recent 11 months, 3 weeks ago

Selected Answer: A

automated failover across AWS Region + minimize latency -> Global Accelerator

automated failover across AWS Region + minimize latency -> Global Accelerator
upvoted 2 times

Murtadhaceit 1 year, 5 months ago

Selected Answer: A

VoIP ==> UDP ==> Global Accelerator.
upvoted 3 times

kaleemanjum 1 year, 6 months ago

Selected Answer: A

AWS Global Accelerator: AWS Global Accelerator is a service that uses static IP addresses (Anycast IPs) to provide a global entry point for your applications. It routes traffic over the AWS global network to the optimal AWS endpoint based on health, geography, and routing policies.

Health Checks: AWS Global Accelerator supports health checks, allowing it to route traffic only to healthy endpoints. This helps in achieving high availability and automated failover across AWS Regions.
upvoted 2 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

A
https://aws.amazon.com/global-accelerator/faqs/#:~:text=Global%20Accelerator%20is%20a%20good,AWS%20Shield%20for%20DDoS%20protection.
upvoted 2 times

ekisako 1 year, 7 months ago

Selected Answer: A

https://docs.aws.amazon.com/global-accelerator/latest/dg/introduction-benefits-of-migrating.html
upvoted 3 times

cciesam 1 year, 7 months ago

Selected Answer: A

Global Accelerator is the answer as it can handle both TCP and UDP
upvoted 2 times

Sugarbear_01 1 year, 7 months ago

Selected Answer: C

This answer should be C
upvoted 1 times

pentium75 1 year, 5 months ago

CloudFront is not for VoIP (which usually uses UDP).
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #648 *Topic 1*

A weather forecasting company needs to process hundreds of gigabytes of data with sub-millisecond latency. The company has a high performance computing (HPC) environment in its data center and wants to expand its forecasting capabilities.

A solutions architect must identify a highly available cloud storage solution that can handle large amounts of sustained throughput. Files that are stored in the solution should be accessible to thousands of compute instances that will simultaneously access and process the entire dataset.

What should the solutions architect do to meet these requirements?

A. Use Amazon FSx for Lustre scratch file systems.

B. Use Amazon FSx for Lustre persistent file systems. **Most Voted**

C. Use Amazon Elastic File System (Amazon EFS) with Bursting Throughput mode.

D. Use Amazon Elastic File System (Amazon EFS) with Provisioned Throughput mode.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

potomac Highly Voted 1 year, 7 months ago

Selected Answer: B

Option A (Amazon FSx for Lustre scratch file systems) is designed for temporary data storage and does not provide the data persistence required for this scenario.

upvoted 9 times

FlyingHawk Most Recent 4 months, 1 week ago

Selected Answer: B

Amazon FSx for Lustre is a fully managed file system optimized for high-performance computing (HPC) workloads.

It provides sub-millisecond latency and high throughput, making it ideal for processing large datasets like weather forecasting data.

It supports thousands of compute instances accessing the file system simultaneously, ensuring scalability for the company's HPC

It supports thousands of compute instances accessing the file system simultaneously, ensuring scalability for the company's HPC environment.

upvoted 1 times

FlyingHawk 4 months, 1 week ago

A persistent file system is recommended because it provides high durability and availability.

Data is replicated within an Availability Zone (AZ) and automatically backed up to Amazon S3, ensuring data is not lost in case of failures.

This is critical for a weather forecasting company that needs to ensure data integrity and availability.

upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: B

HPC + the entire dataset -> FSx Lustre presistence

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

https://docs.aws.amazon.com/fsx/latest/LustreGuide/using-fsx-lustre.html

Both AB can handle the processing requirements but B is Highly Available which is also a requirement not met by A.

CD won't mee the performance requirements

upvoted 4 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

high performance computing, highly available cloud storage solution = Amazon FSx for Lustre persistent file systems

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #649 *Topic 1*

An ecommerce company runs a PostgreSQL database on premises. The database stores data by using high IOPS Amazon Elastic Block Store (Amazon EBS) block storage. The daily peak I/O transactions per second do not exceed 15,000 IOPS. The company wants to migrate the database to Amazon RDS for PostgreSQL and provision disk IOPS performance independent of disk storage capacity.

Which solution will meet these requirements MOST cost-effectively?

A. Configure the General Purpose SSD (gp2) EBS volume storage type and provision 15,000 IOPS.

B. Configure the Provisioned IOPS SSD (io1) EBS volume storage type and provision 15,000 IOPS.

C. Configure the General Purpose SSD (gp3) EBS volume storage type and provision 15,000 IOPS. **Most Voted**

D. Configure the EBS magnetic volume type to achieve maximum IOPS.

Correct Answer: *C*

Community vote distribution

C (100%)

Comments

BillaRanga Highly Voted 10 months ago

GP2 - • Size of the volume and IOPS are linked, max IOPS is 16,000
GP3 - Can increase IOPS up to 16,000 and throughput up to 1000 MiB/s independently

GP3 is 20% cheaper than GP2

upvoted 6 times

Oblako Highly Voted 1 year ago

Selected Answer: C

Both gp2 and gp3 can provision up to 16.000 IOPS. gp3 is cheaper than gp2.

upvoted 5 times

TariqKipkemei Most Recent 1 year ago

Selected Answer: C

MOST cost-effective =GP3

upvoted 4 times

SHAAHIBHUSHANAWS 1 year ago

C
https://aws.amazon.com/ebs/general-purpose/
upvoted 3 times

lagorb 1 year ago

gp2 and gp3 can provision up to 16.000 IOPS, and gp3 is cheaper than gp2
upvoted 3 times

potomac 1 year, 1 month ago

Selected Answer: C

GP3 is better and cheaper than GP2
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #650 *Topic 1*

A company wants to migrate its on-premises Microsoft SQL Server Enterprise edition database to AWS. The company's online application uses the database to process transactions. The data analysis team uses the same production database to run reports for analytical processing. The company wants to reduce operational overhead by moving to managed services wherever possible.

Which solution will meet these requirements with the LEAST operational overhead?

A. Migrate to Amazon RDS for Microsoft SOL Server. Use read replicas for reporting purposes **Most Voted**

B. Migrate to Microsoft SQL Server on Amazon EC2. Use Always On read replicas for reporting purposes

C. Migrate to Amazon DynamoDB. Use DynamoDB on-demand replicas for reporting purposes

D. Migrate to Amazon Aurora MySQL. Use Aurora read replicas for reporting purposes

Correct Answer: *A*

Community vote distribution

A (100%)

Comments

BillaRanga Highly Voted 10 months ago

Selected Answer: A

B - Not the LEAST operational Overhead.
C - It is No-Sql - Not compatible with SQL server which is SQL
D - MS Sql Server to MySQL may miss out some SQL Server functionalities.

A - Read replicas for RDS is easy to create and also it is Asynchronous which should not be a problem for the analytics teams as they can bear 2-3 minutes delay

upvoted 7 times

superalaga Highly Voted 11 months, 3 weeks ago

Selected Answer: A

You can migrate with both A&B but option A is LEAST operational overhead/
A: https://aws.amazon.com/tutorials/move-to-managed/migrate-sql-server-to-amazon-rds/
B: https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/migrate-a-microsoft-sql-server-database-to-aurora-mysql-by-using-aws-dms-and-aws-sct.html

upvoted 6 times

Firdous586 Most Recent 10 months, 3 weeks ago

A is the correct answer since RDS supports OLAP
And aurora OLTP

upvoted 5 times

TariqKipkemei 1 year ago

Selected Answer: A

Only Amazon RDS allows the creation of readable standby DB instances.

upvoted 4 times

potomac 1 year, 1 month ago

Selected Answer: A

A is the only choice

upvoted 6 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #651 *Topic 1*

A company stores a large volume of image files in an Amazon S3 bucket. The images need to be readily available for the first 180 days. The images are infrequently accessed for the next 180 days. After 360 days, the images need to be archived but must be available instantly upon request. After 5 years, only auditors can access the images. The auditors must be able to retrieve the images within 12 hours. The images cannot be lost during this process.

A developer will use S3 Standard storage for the first 180 days. The developer needs to configure an S3 Lifecycle rule.

Which solution will meet these requirements MOST cost-effectively?

A. Transition the objects to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 180 days. S3 Glacier Instant Retrieval after 360 days, and S3 Glacier Deep Archive after 5 years.

B. Transition the objects to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 180 days. S3 Glacier Flexible Retrieval after 360 days, and S3 Glacier Deep Archive after 5 years.

C. Transition the objects to S3 Standard-Infrequent Access (S3 Standard-IA) after 180 days, S3 Glacier Instant Retrieval after 360 days, and S3 Glacier Deep Archive after 5 years. Most Voted

D. Transition the objects to S3 Standard-Infrequent Access (S3 Standard-IA) after 180 days, S3 Glacier Flexible Retrieval after 360 days, and S3 Glacier Deep Archive after 5 years.

Correct Answer: *C*

Community vote distribution

C (85%) Other (15%)

Comments

TariqKipkemei Highly Voted 1 year ago

Selected Answer: C

Images cannot be lost = high availability.
Transition the objects to S3 Standard-Infrequent Access (S3 Standard-IA) after 180 days, S3 Glacier Instant Retrieval after 360 days, and S3 Glacier Deep Archive after 5 years.

upvoted 11 times

dilaaziz Highly Voted 1 year, 1 month ago

Selected Answer: C

https://aws.amazon.com/s3/storage-classes/glacier/
upvoted 6 times

Salilgen Most Recent 5 months, 1 week ago

Selected Answer: C

IMO answer is C
upvoted 1 times

Linuslin 6 months, 3 weeks ago

Selected Answer: C

"The developer needs to configure an S3 Lifecycle rule."--->One Zone-IA can't transfer to Glacier Instant Retrieval--->A is out.
Check - Unsupported lifecycle transitions
https://docs.aws.amazon.com/AmazonS3/latest/userguide/lifecycle-transition-general-considerations.html
"Images cannot be lost = high availability"--->Can't be One Zone-IA--->B is out.
"the images need to be archived but must be available instantly upon request"--->Can't be "Flexible" Retrieval--->D is out.

Only C is the correct answer.
upvoted 3 times

Neung983 9 months, 1 week ago

Selected Answer: A

A.
Here's why this option is the most cost-effective:
+S3 One Zone-IA (after 180 days): Offers lower storage costs compared to S3 Standard for infrequently accessed data (180 - 360 days) while maintaining good availability for retrieval.
+S3 Glacier Instant Retrieval (after 360 days): Provides immediate access to archived images (360 - 5 years) at a significantly lower cost than S3 Standard storage. Retrieval costs are incurred but typically lower than keeping the data in S3 Standard.
+S3 Glacier Deep Archive (after 5 years): Offers the lowest storage cost for long-term archival (beyond 5 years) with retrieval times within 12 hours, meeting the auditor access requirement and minimizing ongoing storage costs.
upvoted 1 times

Salilgen 5 months, 1 week ago

A is wrong because "The images cannot be lost during this process" then you can use OneZone
upvoted 1 times

Antitouch 11 months, 1 week ago

Selected Answer: B

https://aws.amazon.com/s3/storage-classes/glacier/#:~:text=S3%20Glacier%20Flexible%20Retrieval%20delivers,year%20and%20is%20retrieved%20asynchronously.
S3 Glacier Flexible Retrieval delivers low-cost storage, up to 10% lower cost than S3 Glacier Instant Retrieval. Flexible retrieval is cheaper than Instant retrieval.
S3 Glacier Flexible retrieval storage class provides minutes to 12 hours retrieval of data. Which is within the required time by auditors.
--> We should select flexible retrieval.

The design is not caring about the high availability. The design is caring about cost. One zone-IA is cheaper than standard IA.
--> We should select One Zone IA.
upvoted 1 times

awsgeek75 11 months ago

"The images cannot be lost during this process" is a core requirement.
The design cares about data loss and 5 years is a long time and AZ failure will result in data loss.
upvoted 4 times

pentium75 11 months, 1 week ago

Selected Answer: C

A, B impose risk of the images being lost in case of AZ failure
D does not allow instant access after 180 days
upvoted 4 times

ale_brd_111 11 months, 2 weeks ago

Selected Answer: C

Images cannot be lost = high availability. A exposes images to risk
upvoted 2 times

Alex1atd 1 year ago

Selected Answer: C

The images cannot be lost during this process.
upvoted 3 times

1rob 1 year ago

Selected Answer: C

"The images cannot be lost during this process" , imho this rules out S3 One zone infrequent access. S3 Glacier Instant Retrieval gives immediate access. S3 Glacier Flexible Retrieval does not give immediate access. so C.
upvoted 5 times

EdenWang 1 year ago

Selected Answer: A

high availability is not mentioned, thus I go for A
upvoted 1 times

pentium75 11 months, 1 week ago

"The images cannot be lost during this process."
upvoted 5 times

TheLaPlanta 8 months, 3 weeks ago

That's not HA
upvoted 1 times

cciesam 1 year, 1 month ago

Selected Answer: A

I'll go for A as it doesn't talk about High availability. Considering cost. I'll go for A
upvoted 3 times

ekisako 1 year, 1 month ago

"The images cannot be lost during this process."
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #652 *Topic 1*

A company has a large data workload that runs for 6 hours each day. The company cannot lose any data while the process is running. A solutions architect is designing an Amazon EMR cluster configuration to support this critical data workload.

Which solution will meet these requirements MOST cost-effectively?

A. Configure a long-running cluster that runs the primary node and core nodes on On-Demand Instances and the task nodes on Spot Instances.

B. Configure a transient cluster that runs the primary node and core nodes on On-Demand Instances and the task nodes on Spot Instances. **Most Voted**

C. Configure a transient cluster that runs the primary node on an On-Demand Instance and the core nodes and task nodes on Spot Instances.

D. Configure a long-running cluster that runs the primary node on an On-Demand Instance, the core nodes on Spot Instances, and the task nodes on Spot Instances.

Correct Answer: *B*

Community vote distribution

B (100%)

Comments

louisaok Highly Voted 1 year, 7 months ago

Relax man. take a break since you have made this far so far.

upvoted 54 times

AWSSURI 9 months, 1 week ago

Rest at the end NOT in the middle
-- Kobe Bryant

upvoted 9 times

potomac Highly Voted 1 year, 7 months ago

Selected Answer: B

A transient cluster provides cost savings because it runs only during the computation time, and it provides scalability and flexibility in a cloud environment flexibility in a cloud environment.

Option C (transient cluster with On-Demand primary node and Spot core and task nodes) exposes the core nodes to Spot Instance interruptions, which may not be acceptable for a workload that cannot lose any data.

upvoted 14 times

EllenLiu Most Recent 5 months, 2 weeks ago

Selected Answer: B

take a break and feel better now. thanks for encouraging with each other!
daily batch handling within several hours => transient cluster
data-critical application => spot for task node only

upvoted 1 times

awsgeek75 1 year, 4 months ago

Selected Answer: B

AD are long-running so don't fit in with 6 hours schedule
BC are ideal for scheduled EMR activities
C is wrong as running core node on Spot instance has a risk of data loss
https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-master-core-task-nodes.html
B is correct because primary, core will be stable on on-demand as recommended by AWS and task can go on spot instances as task nodes are short lived by nature anyway

upvoted 6 times

pentium75 1 year, 5 months ago

Selected Answer: B

"Long-running cluster" = runs until you shut it down
"Transient cluster" = runs until the workload is completed
This runs only 6 hours each day -> transient -> B or C

"Cannot lose any data while the process is running" -> Primary and core nodes cannot be Spot instances -> A or B
https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-plan-longrunning-transient.html
https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-plan-instances-guidelines.html

upvoted 10 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

Cannot loose data = ondemand primary + core nodes
Save on costs = spot task nodes
Runs for 6 hours = transient cluster

upvoted 6 times

SHAAHIBHUSHANAWS 1 year, 6 months ago

A
https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-plan-instances-guidelines.html
It's long running and no data loss is needed.

upvoted 2 times

pentium75 1 year, 5 months ago

The link says you can lose data if you are running a transient cluster WITH ONLY Spot instances. "Long-running" = runs until you shut it down, "Transient" = Runs until the workload is completed

upvoted 1 times

whiterick 1 year, 4 months ago

Option A suggests a long-running cluster, which continues to run until manually terminated. This means that even if tasks are rerouted due to Spot Instance interruptions, the cluster itself remains active, allowing the rerouted tasks to complete on other nodes.

Option B suggests a transient cluster, which is terminated after all steps are completed. If the Spot Instances are interrupted and tasks are not completed, the cluster might still terminate after the steps are deemed complete, potentially leading to incomplete processing of data.

upvoted 1 times

MFKang 1 year, 6 months ago

Get up Stand up
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #653 *Topic 1*

A company maintains an Amazon RDS database that maps users to cost centers. The company has accounts in an organization in AWS Organizations. The company needs a solution that will tag all resources that are created in a specific AWS account in the organization. The solution must tag each resource with the cost center ID of the user who created the resource.

Which solution will meet these requirements?

A. Move the specific AWS account to a new organizational unit (OU) in Organizations from the management account. Create a service control policy (SCP) that requires all existing resources to have the correct cost center tag before the resources are created. Apply the SCP to the new OU.

B. Create an AWS Lambda function to tag the resources after the Lambda function looks up the appropriate cost center from the RDS database. Configure an Amazon EventBridge rule that reacts to AWS CloudTrail events to invoke the Lambda function. **Most Voted**

C. Create an AWS CloudFormation stack to deploy an AWS Lambda function. Configure the Lambda function to look up the appropriate cost center from the RDS database and to tag resources. Create an Amazon EventBridge scheduled rule to invoke the CloudFormation stack.

D. Create an AWS Lambda function to tag the resources with a default value. Configure an Amazon EventBridge rule that reacts to AWS CloudTrail events to invoke the Lambda function when a resource is missing the cost center tag.

Correct Answer: *B*

Community vote distribution

B (64%) A (36%)

Comments

1Alpha1 Highly Voted 1 year, 4 months ago

Selected Answer: A

I'm not sure, but I think this question is from professional solution architect question pool.
Please have a look at this one as well.
https://www.examtopics.com/discussions/amazon/view/112780-exam-aws-certified-solutions-architect-professional-sap-c02/
upvoted 10 times

JohnYu 7 months, 3 weeks ago

SCPs cannot directly apply tags to resources; they can only restrict actions based on policies. They do not automatically tag resources. SCPs are also used to enforce policies across accounts but cannot look up values from an RDS database or dynamically assign tags.
upvoted 3 times

pentium75 Highly Voted 1 year, 5 months ago

Selected Answer: B

A policy cannot look up "the cost center ID of the user who created the resource", we need Lambda to do that. Thus A is out.

C would work but runs on a schedule which doesn't make sense (and we would temporarily have untagged resources).

D tags resources "with a default value" which is not what we want.
upvoted 8 times

Ernestokoro 1 year, 5 months ago

Please how do you account for this part of the question with option B "The solution must tag each resource with the cost center ID of the user who created the resource." ? For me this typically what SCP would handle.
upvoted 1 times

awsgeek75 1 year, 4 months ago

That SCP won't know the cost centre. "RDS database that maps users to cost centers"
Unless the solution can read the RDS, it won't work and SCP cannot be programmed to read from RDS before applying the cost centre.
upvoted 3 times

FlyingHawk Most Recent 4 months, 2 weeks ago

Selected Answer: B

A is incorrect : SCPs are used to enforce permissions and restrictions across accounts in an organization. They cannot be used to dynamically tag resources based on user information stored in an RDS database. This option does not meet the requirement to tag resources with the cost center ID of the user who created them. C incorrect as it only deploy the CF stack with Lambda, but no mention how lambda function will be triggered. D default value is incorrect.
upvoted 2 times

bujuman 7 months ago

Selected Answer: B

SCP can not be used to Tag resources
upvoted 2 times

tonybuivannghia 7 months, 2 weeks ago

Selected Answer: B

B is my choice
upvoted 2 times

1e22522 10 months ago

Selected Answer: B

The best solution that meets all the requirements is option B. Here's why:

It can tag all resources created in a specific AWS account within the organization.
It uses a Lambda function to look up the appropriate cost center from the RDS database, ensuring accurate tagging.
The EventBridge rule reacting to CloudTrail events ensures that resources are tagged as they are created.
This approach can dynamically tag each resource with the cost center ID of the user who created it.
upvoted 5 times

sheilawu 1 year ago

Selected Answer: A

I will vote for A
upvoted 1 times

sandordini 1 year, 1 month ago

Selected Answer: B

We need a solution, to automatically tag, also the existing resources. A,C, are more or less working solutions for new resources, but neither can do the tagging of existing resources. D would add a default tag instead of the specific CC.
upvoted 2 times

gsgdga 1 year, 2 months ago

Selected Answer: A

A is right
https://docs.aws.amazon.com/ko_kr/organizations/latest/userguide/orgs_tagging_abac.html
upvoted 4 times

fea9bdf 1 year, 5 months ago

Answer is A, SCP handles the assignment, no need for a Lambda function, that's unnecessary t seems like
Service Control Policies (SCPs)

SCPs are a policy type that you can utilize to manage permissions across accounts in your AWS Organization.
Using SCPs lets you ensure that your accounts stay within your organization’s access control guidelines.
SCPs can be used along-side tag policies to ensure that the tags are applied at the resource creation time and remain attached to the resource.
upvoted 1 times

pentium75 1 year, 5 months ago

How would an SCP look up the correct cost center in the database?
upvoted 4 times

ale_brd_111 1 year, 5 months ago

Selected Answer: B

the company still maintains the RDS, nowhere was asked to drop using it, therefore we shall use a solution that takes advantages of it.
upvoted 4 times

ftaws 1 year, 5 months ago

Selected Answer: A

I also choose A.
upvoted 2 times

awsgeek75 1 year, 4 months ago

SCP cannot connect to RDS where the cost centre information is stored so A won't work.
upvoted 2 times

Oluwatosin09 1 year, 1 month ago

Awsgeek75 and Pentium must be the same person.
They always have the same answers with contributions always on point.
Great Job!
upvoted 2 times

Cyberkayu 1 year, 5 months ago

Selected Answer: A

Company have Organization. A specific AWS account need to ensure all resources were tagged.

Move this specific AWS account under the company OU, use SCP to enforce top down policies that every member account to adhere.

Answer A.
upvoted 1 times

pentium75 1 year, 5 months ago

How would an SCP look up the correct cost center in the database?
upvoted 2 times

upvoted 2 times

evelynsun 1 year, 5 months ago

Selected Answer: B

sorry, i would choose B.
because it allows you to tag resources as they are created, without requiring you to move existing resources.

upvoted 2 times

evelynsun 1 year, 5 months ago

Selected Answer: A

This solution is the best way to meet the requirements of the company. It ensures that all resources in the specific AWS account are tagged with the cost center ID of the user who created the resource. It also allows the company to easily manage and enforce compliance with its tagging policies.

upvoted 1 times

pentium75 1 year, 5 months ago

How would an SCP look up the correct cost center in the database?

upvoted 2 times

TariqKipkemei 1 year, 6 months ago

Selected Answer: B

Create an AWS Lambda function to tag the resources after the Lambda function looks up the appropriate cost center from the RDS database. Configure an Amazon EventBridge rule that reacts to AWS CloudTrail events to invoke the Lambda function.

upvoted 2 times

t0nx 1 year, 6 months ago

Selected Answer: B

This solution utilizes AWS Lambda and Amazon EventBridge to automate the tagging process based on information from the RDS database and CloudTrail events.

AWS Lambda Function: Create a Lambda function that can look up the cost center information from the RDS database and tag resources accordingly.

Amazon EventBridge Rule: Set up an EventBridge rule to react to AWS CloudTrail events. The rule triggers the Lambda function whenever a resource is created, allowing dynamic tagging based on the cost center associated with the user in the RDS database.

This solution provides automation, ensuring that resources are tagged appropriately with the cost center ID of the user who created the resource. It also allows for flexibility in updating cost center information without modifying the infrastructure.

upvoted 5 times

EXAMTOPICS

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #654 *Topic 1*

A company recently migrated its web application to the AWS Cloud. The company uses an Amazon EC2 instance to run multiple processes to host the application. The processes include an Apache web server that serves static content. The Apache web server makes requests to a PHP application that uses a local Redis server for user sessions.

The company wants to redesign the architecture to be highly available and to use AWS managed solutions.

Which solution will meet these requirements?

A. Use AWS Elastic Beanstalk to host the static content and the PHP application. Configure Elastic Beanstalk to deploy its EC2 instance into a public subnet. Assign a public IP address.

B. Use AWS Lambda to host the static content and the PHP application. Use an Amazon API Gateway REST API to proxy requests to the Lambda function. Set the API Gateway CORS configuration to respond to the domain name. Configure Amazon ElastiCache for Redis to handle session information.

C. Keep the backend code on the EC2 instance. Create an Amazon ElastiCache for Redis cluster that has Multi-AZ enabled. Configure the ElastiCache for Redis cluster in cluster mode. Copy the frontend resources to Amazon S3. Configure the backend code to reference the EC2 instance.

D. Configure an Amazon CloudFront distribution with an Amazon S3 endpoint to an S3 bucket that is configured to host the static content. Configure an Application Load Balancer that targets an Amazon Elastic Container Service (Amazon ECS) service that runs AWS Fargate tasks for the PHP application. Configure the PHP application to use an Amazon ElastiCache for Redis cluster that runs in multiple Availability Zones. **Most Voted**

Correct Answer: *D*

Community vote distribution

D (100%)

Comments

awsgeek75 Highly Voted 11 months ago

Selected Answer: D

Key requirements: HA and Managed Services
Key components: PHP, Static content, Redis Elasticache
AB are instantly useless for static content scaling