

Kinesis Data Streams: This is designed for real-time data ingestion, which is exactly what the scenario requires.

Kinesis Data Firehose: This service can then be used to efficiently deliver the ingested data from Kinesis streams to an S3 bucket, fulfilling the need to store raw data in S3

upvoted 1 times

JA2018 6 months, 1 week ago

Why the other options are incorrect:

B. AWS Database Migration Service (DMS):

DMS is primarily used for migrating databases, not for real-time data ingestion from third-party applications.

C. AWS DataSync:

While DataSync can transfer data to S3, it's not optimized for real-time data ingestion and would likely not be the best choice for this scenario.

D. AWS Direct Connect:

Direct Connect is used for dedicated private network connections between your on-premises network and AWS, not for real-time data ingestion from third-party applications.

upvoted 1 times

JA2018 6 months, 1 week ago

Key points to remember:

When dealing with real-time data ingestion, Kinesis Data Streams is generally the preferred solution on AWS.

Kinesis Data Firehose can be used to easily stream data from Kinesis to various destinations like S3 buckets.

upvoted 1 times

mk168898 7 months, 3 weeks ago

Selected Answer: A

real time data => amazon kinesis data stream

upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: A

- Amazon Kinesis Data Streams: designed for real-time data ingestion.

- Kinesis Data Firehose: can consume data from Kinesis Data Streams and automatically deliver it to Amazon S3.

A is the answer.

upvoted 2 times

Mayank0502 11 months, 1 week ago

Selected Answer: C

each ec2 needs to proceed data separately

upvoted 1 times

xBUGx 1 year, 2 months ago

Selected Answer: C

A is best solution, but i think the question is saying "The application needs to ingest real-time data from third-party applications." and the application is run on EC2.

so i think we need a solution that works with the application on EC2 for this question?

upvoted 3 times

Sergiuss95 1 year, 1 month ago

DataSync is more suitable for transferring data between on-premises storage systems and AWS, rather than ingesting real-time data. Best solution is A

upvoted 5 times

seetpt 1 year, 3 months ago

Agree with A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #819

Topic 1

A company's application is receiving data from multiple data sources. The size of the data varies and is expected to increase over time. The current maximum size is 700 KB. The data volume and data size continue to grow as more data sources are added.

The company decides to use Amazon DynamoDB as the primary database for the application. A solutions architect needs to identify a solution that handles the large data sizes.

Which solution will meet these requirements in the MOST operationally efficient way?

- A. Create an AWS Lambda function to filter the data that exceeds DynamoDB item size limits. Store the larger data in an Amazon DocumentDB (with MongoDB compatibility) database.
- B. Store the large data as objects in an Amazon S3 bucket. In a DynamoDB table, create an item that has an attribute that points to the S3 URL of the data. **Most Voted**
- C. Split all incoming large data into a collection of items that have the same partition key. Write the data to a DynamoDB table in a single operation by using the BatchWriteItem API operation.
- D. Create an AWS Lambda function that uses gzip compression to compress the large objects as they are written to a DynamoDB table.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**Neung983** **Highly Voted** 1 year, 3 months ago**Selected Answer: B**

option B is the most operationally efficient solution for handling large data sizes in Amazon DynamoDB.
upvoted 9 times

seetpt **Highly Voted** 1 year, 3 months ago**Selected Answer: B**

B is correct

upvoted 5 times

LeonSauveterre Most Recent 5 months, 2 weeks ago

Selected Answer: B

A - Introducing another database. No need.

B - Decoupling the large data storage (S3) from the metadata or smaller structured data (DynamoDB) is an AWS-recommended pattern.

C - OK but this is too complicated.

D - Compression does not guarantee that the data will always fit within DynamoDB's 400 KB limit.

upvoted 2 times

Scheldon 12 months ago

Selected Answer: B

AnswerB

Compression of data in DynamoDB is a good idea especially for text data link from forum, but to do that we do not need AWS Lambda if I'm not wrong.

In other hand Storing big object on S3 and saving URL to it in DynamoDB is one of best practices mentioned by Amazon.

Hence we do not know what kind of data we are storing in DB and how big objects will be in the future option B looks like the best solution.

<https://aws.amazon.com/blogs/database/large-object-storage-strategies-for-amazon-dynamodb/> <<<< Read Option 2

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/bp-use-s3-too.html>

upvoted 1 times

Sergiuss95 1 year, 1 month ago

Selected Answer: B

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/bp-use-s3-too.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #820

Topic 1

A company is migrating a legacy application from an on-premises data center to AWS. The application relies on hundreds of cron jobs that run between 1 and 20 minutes on different recurring schedules throughout the day.

The company wants a solution to schedule and run the cron jobs on AWS with minimal refactoring. The solution must support running the cron jobs in response to an event in the future.

Which solution will meet these requirements?

- A. Create a container image for the cron jobs. Use Amazon EventBridge Scheduler to create a recurring schedule. Run the cron job tasks as AWS Lambda functions.
- B. Create a container image for the cron jobs. Use AWS Batch on Amazon Elastic Container Service (Amazon ECS) with a scheduling policy to run the cron jobs.
- C. Create a container image for the cron jobs. Use Amazon EventBridge Scheduler to create a recurring schedule. Run the cron job tasks on AWS Fargate. **Most Voted**
- D. Create a container image for the cron jobs. Create a workflow in AWS Step Functions that uses a Wait state to run the cron jobs at a specified time. Use the RunTask action to run the cron job tasks on AWS Fargate.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Kezuko **Highly Voted** 8 months, 3 weeks ago

Give yourself a pat on the back when you reach this question, its been a long run
upvoted 19 times

Drew3000 8 months, 2 weeks ago

I finally managed to get through the last question, then refreshed the page , and they have added more questions.
upvoted 7 times

LeonSauveterre 5 months, 2 weeks ago

I know right. Where did they get so many questions? It's been a long run but too long now...
upvoted 1 times

Linuslin Highly Voted 6 months, 2 weeks ago

Selected Answer: C

Lambda has 15 mins limit, so A is out.

B works, but you have to run highly-available virtual machines or containers waiting for the event to happen.

C is the best answer in this question, with AWS Fargate allows you to pay for only what you use and free you from provisioning, configuring, and scaling clusters of Amazon EC2 instances.

<https://aws.amazon.com/blogs/containers/migrate-cron-jobs-to-event-driven-architectures-using-amazon-elastic-container-service-and-amazon-eventbridge/>

upvoted 5 times

FlyingHawk Most Recent 4 months, 2 weeks ago

Selected Answer: C

As an example, one customer had hundreds of cron jobs configured on a single on-premise server. They chose to run these various cron jobs as separate scheduled tasks on Amazon Elastic Container Service (Amazon ECS) using AWS Fargate. They used Amazon EventBridge to schedule the tasks because it triggers an Amazon ECS task at scheduled times rather than having to run highly-available virtual machines or containers waiting for the event to happen. They chose Amazon ECS because of its simple container orchestration. And they chose AWS Fargate because it allows them to pay for only what they use and frees them from provisioning, configuring, and scaling clusters of Amazon EC2 instances.

upvoted 1 times

cvoiceip 9 months ago

Ans : C

<https://aws.amazon.com/blogs/containers/migrate-cron-jobs-to-event-driven-architectures-using-amazon-elastic-container-service-and-amazon-eventbridge/>

upvoted 1 times

seetpt 9 months ago

Selected Answer: C

C because lambda has 15min time limit.

upvoted 2 times

asdfcdsxdc 9 months ago

Selected Answer: C

its either A or C. C looks correct because lambda works for 15 mins and the question says between 1-20

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #821

Topic 1

A company uses Salesforce. The company needs to load existing data and ongoing data changes from Salesforce to Amazon Redshift for analysis. The company does not want the data to travel over the public internet.

Which solution will meet these requirements with the LEAST development effort?

- A. Establish a VPN connection from the VPC to Salesforce. Use AWS Glue DataBrew to transfer data.
- B. Establish an AWS Direct Connect connection from the VPC to Salesforce. Use AWS Glue DataBrew to transfer data.
- C. Create an AWS PrivateLink connection in the VPC to Salesforce. Use Amazon AppFlow to transfer data. **Most Voted**
- D. Create a VPC peering connection to Salesforce. Use Amazon AppFlow to transfer data.

Correct Answer: C

Community vote distribution

C (100%)

Comments

LeonSauveterre 5 months, 2 weeks ago

Selected Answer: C

- A - VPN connection, public internet, no.
 - B - Significant setup effort and cost. Also, DataBrew is not suitable for loading ongoing data changes.
 - C - Amazon AppFlow is a fully managed integration service that natively supports Salesforce and Amazon Redshift.
 - D - VPC peering is for connecting two AWS VPCs. Salesforce is external to AWS, so no.
- upvoted 3 times

LeonSauveterre 5 months, 2 weeks ago

Also, maybe this will help you along the way (things after "==" are their disadvantages):

1. VPN Connection: For secure, encrypted traffic between on-premises and AWS over the public internet. ==> Relies on public internet (higher latency)
2. Direct Connect: For dedicated private network connectivity between on-premises and AWS (like data centers with large workloads). ==> Higher cost, longer setup time.
3. PrivateLink: For secure private connectivity to AWS services or third-party services (like Salesforce) within AWS. ==> Simple setup but limited to services configured with PrivateLink (not general connectivity).

4. VPC Peering: For direct connectivity between two AWS VPCs (like multi-account or region scenarios). ==> Does not support transitive connections; not suitable for external services like Salesforce.

upvoted 2 times

mk168898 7 months, 3 weeks ago

3rd party SaaS salesforce integration => Use AWS AppFlow

So left C and D

Not D because VPC Peering need 2 VPC and 3rd party SaaS does not have a VPC

upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: C

AWS PrivateLink: This service enables private connectivity between VPCs and supported AWS services, effectively keeping data off the public internet. It allows secure communication without exposing your data to the internet.

Amazon AppFlow: This is a fully managed integration service that simplifies data transfer between SaaS applications (like Salesforce) and AWS services (like Amazon Redshift).

upvoted 2 times

Johnoppong101 9 months, 3 weeks ago

Selected Answer: C

C is the answer

upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: C

Private link for sure

upvoted 2 times

Scheldon 12 months ago

Selected Answer: C

AnswerC

To connect your own VPC with third-party VPC we need to use PrivateLink.

AWS PrivateLink is a highly available, scalable technology that you can use to privately connect your VPC to services as if they were in your VPC. You do not need to use an internet gateway, NAT device, public IP address, AWS Direct Connect connection, or AWS Site-to-Site VPN connection to allow communication with the service from your private subnets. Therefore, you control the specific API endpoints, sites, and services that are reachable from your VPC.

<https://docs.aws.amazon.com/vpc/latest/privatelink/what-is-privatelink.html>

upvoted 4 times

Kaula 1 year, 2 months ago

Selected Answer: C

Should be C

upvoted 2 times

Kaula 1 year, 2 months ago

C

<https://docs.aws.amazon.com/connect/latest/adminguide/integrate-salesforce-tasks.html>

<https://docs.aws.amazon.com/connect/latest/adminguide/vpc-interface-endpoints.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #822

Topic 1

A company recently migrated its application to AWS. The application runs on Amazon EC2 Linux instances in an Auto Scaling group across multiple Availability Zones. The application stores data in an Amazon Elastic File System (Amazon EFS) file system that uses EFS Standard-Infrequent Access storage. The application indexes the company's files. The index is stored in an Amazon RDS database.

The company needs to optimize storage costs with some application and services changes.

Which solution will meet these requirements MOST cost-effectively?

- A. Create an Amazon S3 bucket that uses an Intelligent-Tiering lifecycle policy. Copy all files to the S3 bucket. Update the application to use Amazon S3 API to store and retrieve files. **Most Voted**
- B. Deploy Amazon FSx for Windows File Server file shares. Update the application to use CIFS protocol to store and retrieve files.
- C. Deploy Amazon FSx for OpenZFS file system shares. Update the application to use the new mount point to store and retrieve files.
- D. Create an Amazon S3 bucket that uses S3 Glacier Flexible Retrieval. Copy all files to the S3 bucket. Update the application to use Amazon S3 API to store and retrieve files as standard retrievals.

Correct Answer: A

Community vote distribution

A (100%)

Comments

xBUGx **Highly Voted** 1 year, 2 months ago

Selected Answer: A

i go with A since there is no other better options
upvoted 6 times

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: A

B, C, D off.

B - Windows

C - OpenZFS Costs ~ NFS Costs *3 (optimized mainly for high-performance Data analysis)

D - Glacier Standard retrieval: 12 Hrs

upvoted 5 times

Dantecito 3 months, 2 weeks ago

Expedited retrievals: Typically complete in 1–5 minutes

Standard retrievals: Typically complete in 3–5 hours

Free bulk retrievals: Typically complete in 5–12 hours

upvoted 1 times

TeamACT 10 months, 3 weeks ago

Glacier standard retrieval is 5 hrs and not 12. It is glacier deep archive standard that is 12 hours.

upvoted 2 times

mk168898 Most Recent 7 months, 3 weeks ago

Selected Answer: A

Chose A because optimise storage costs => s3 bucket intelligent tiering

upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: A

Since the company is ok with "some application and services changes", then A is definitely the most cost-effective option.

D can take up to few hours to complete retrievals.

upvoted 2 times

Johnoppong101 9 months, 3 weeks ago

Selected Answer: A

EFS Infrequent access -> milliseconds retrieval time, can't replace with 12hrs for Glacier.

upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: A

optimize storage costs with some application and services changes -> Intelligent Tiering

upvoted 2 times

Scheldon 12 months ago

Selected Answer: A

AnswerA

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/intelligent-tiering-overview.html>

upvoted 3 times

TruthWS 1 year, 2 months ago

A is correct

upvoted 2 times

Kenneth99 1 year, 2 months ago

should be A?

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #823

Topic 1

A robotics company is designing a solution for medical surgery. The robots will use advanced sensors, cameras, and AI algorithms to perceive their environment and to complete surgeries.

The company needs a public load balancer in the AWS Cloud that will ensure seamless communication with backend services. The load balancer must be capable of routing traffic based on the query strings to different target groups. The traffic must also be encrypted.

Which solution will meet these requirements?

- A. Use a Network Load Balancer with a certificate attached from AWS Certificate Manager (ACM). Use query parameter-based routing.
- B. Use a Gateway Load Balancer. Import a generated certificate in AWS Identity and Access Management (IAM). Attach the certificate to the load balancer. Use HTTP path-based routing.
- C. Use an Application Load Balancer with a certificate attached from AWS Certificate Manager (ACM). Use query parameter-based routing. **Most Voted**
- D. Use a Network Load Balancer. Import a generated certificate in AWS Identity and Access Management (IAM). Attach the certificate to the load balancer. Use query parameter-based routing.

Correct Answer: C*Community vote distribution*C (100%)**Comments****TruthWS** **Highly Voted** 1 year, 2 months ago

Option C - parameter is not a thing NLB can process
upvoted 6 times

LeonSauveterre **Most Recent** 5 months, 1 week ago**Selected Answer: C**

Neither Network Load Balancer nor Gateway Load Balancer supports query parameter-based routing.
upvoted 1 times

mk168898 7 months, 3 weeks ago

Selected Answer: C

only ALB can route traffic based on query parameter, NLB cannot.
So C

upvoted 2 times

Scheldon 12 months ago

Selected Answer: C

AnswerC

I was thinking that it will be network LB but after checks it occurs that only Application LB is able to redirect/forward traffic based on query string. and for encrypted traffic ACM is needed

We recommend that you create certificates for your load balancer using AWS Certificate Manager (ACM). ACM supports RSA certificates with 2048, 3072, and 4096-bit key lengths, and all ECDSA certificates. ACM integrates with Elastic Load Balancing so that you can deploy the certificate on your load balancer. For more information, see the AWS Certificate Manager User Guide.

<https://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-listeners.html#rule-condition-types>

<https://docs.aws.amazon.com/elasticloadbalancing/latest/application/create-https-listener.html>

upvoted 4 times

alawada 1 year, 2 months ago

Selected Answer: C

Provision an Application Load Balancer (ALB) in the AWS Cloud. ALB is a Layer 7 load balancer that supports advanced routing features, including path-based routing.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #824

Topic 1

A company has an application that runs on a single Amazon EC2 instance. The application uses a MySQL database that runs on the same EC2 instance. The company needs a highly available and automatically scalable solution to handle increased traffic.

Which solution will meet these requirements?

- A. Deploy the application to EC2 instances that run in an Auto Scaling group behind an Application Load Balancer. Create an Amazon Redshift cluster that has multiple MySQL-compatible nodes.
- B. Deploy the application to EC2 instances that are configured as a target group behind an Application Load Balancer. Create an Amazon RDS for MySQL cluster that has multiple instances.
- C. Deploy the application to EC2 instances that run in an Auto Scaling group behind an Application Load Balancer. Create an Amazon Aurora Serverless MySQL cluster for the database layer. **Most Voted**
- D. Deploy the application to EC2 instances that are configured as a target group behind an Application Load Balancer. Create an Amazon ElastiCache for Redis cluster that uses the MySQL connector.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**haci** **Highly Voted** 1 year, 2 months ago**Selected Answer: C**

Target groups are just a group of Ec2 instances. Target groups are closely associated with ELB and not ASG. We can just use ELB and Target groups to route requests to EC2 instances. With this setup, there is no autoscaling which means instances cannot be added or removed when your load increases/decreases.

upvoted 8 times

sheilawu **Highly Voted** 1 year ago**Selected Answer: C**

scalable solution= Amazon Aurora Serverless

upvoted 5 times

LeonSauveterre **Most Recent** 5 months, 1 week ago

Selected Answer: C

I don't see how Redshift or ElastiCache relates to the question.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: C

Option B:

Target Group: it doesn't inherently imply automatic scaling. You would need to manage scaling separately, either manually or through other mechanisms like scheduled actions.

Option C:

Auto Scaling Group: This ensures that the EC2 instances can automatically scale in or out based on traffic and demand.

"The company needs a highly available and automatically scalable solution" => C

upvoted 2 times

Scheldon 12 months ago

Selected Answer: C

AnswerC

highly available and automatically scalable solution = Auto Scale for EC2 in front which we will have ALB and Aurora Serverless will give us perfect decoupled solution in which we can increase amount of servers per need and in case of server failure AutoScale will run new EC2 instance

upvoted 3 times

camps 1 year, 2 months ago

It's C!

upvoted 2 times

TruthWS 1 year, 2 months ago

Option C - keywords HA, automatically scalable

upvoted 2 times

alawada 1 year, 2 months ago

Selected Answer: C

C is what I will go for

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #825

Topic 1

A company is planning to migrate data to an Amazon S3 bucket. The data must be encrypted at rest within the S3 bucket. The encryption key must be rotated automatically every year.

Which solution will meet these requirements with the LEAST operational overhead?

A. Migrate the data to the S3 bucket. Use server-side encryption with Amazon S3 managed keys (SSE-S3). Use the built-in key rotation behavior of SSE-S3 encryption keys.


B. Create an AWS Key Management Service (AWS KMS) customer managed key. Enable automatic key rotation. Set the S3 bucket's default encryption behavior to use the customer managed KMS key. Migrate the data to the S3 bucket. **Most Voted**

C. Create an AWS Key Management Service (AWS KMS) customer managed key. Set the S3 bucket's default encryption behavior to use the customer managed KMS key. Migrate the data to the S3 bucket. Manually rotate the KMS key every year.

D. Use customer key material to encrypt the data. Migrate the data to the S3 bucket. Create an AWS Key Management Service (AWS KMS) key without key material. Import the customer key material into the KMS key. Enable automatic key rotation.

Correct Answer: B

Community vote distribution

**Comments**

EllenLiu 5 months, 2 weeks ago

Selected Answer: B

https://repost.aws/questions/QUES_1VN01TU-eRSO3LXergA/s3-managed-key-sse-s3-rotation-period
upvoted 2 times

JoeTromundo 8 months, 2 weeks ago

Selected Answer: B

The answer can't be A. In addition to other justifications written here in the comments, if the data is copied before enabling encryption, this data will not be encrypted.

upvoted 4 times

Johnnopping101 9 months, 3 weeks ago

Selected Answer: B

B is the Answer
upvoted 1 times

n999 10 months, 2 weeks ago

Selected Answer: A

It's said should be encrypted within S3 not before so A its correct
upvoted 2 times

Johnnopping101 9 months, 3 weeks ago

Me: Does SSE-S3 allow custom key rotation scheduling?

Gemini: No, SSE-S3 does not allow for custom key rotation scheduling.

Gemini: If you require more granular control over key rotation, you should consider using Server-Side Encryption with AWS KMS (SSE-KMS)

upvoted 2 times

Scheldon 12 months ago

Selected Answer: B

AnswerB

Looks like key rotation is only possible when KMS is in use. If we will use AWS managed keys Rotation is forced and if we will not provide any specifications regarding rotation time for key, KMS will rotate key every 365days.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingKMSEncryption.html>

<https://docs.aws.amazon.com/kms/latest/developerguide/rotate-keys.html>

<https://docs.aws.amazon.com/kms/latest/developerguide/concepts.html#key-mgmt>

upvoted 3 times

sheilawu 1 year ago

Selected Answer: B

If you see rotation, SEE-SE is out
upvoted 3 times

f07ed8f 1 year ago

Selected Answer: B

SSE-S3 does not rotate the key EVERY YEAR and it is not fit the requirement
upvoted 3 times

Linuslin 1 year ago

Selected Answer: A

This question is flawed.

SSE-S3 is not SSE-KMS, so it will not automatic rotation every year, only KMS will. (check link below)

But the question says "LEAST operational overhead", I think it want us to choose SSE-S3, so I will pick option A.

upvoted 2 times

Linuslin 1 year ago

SSE-S3 is the simplest method to use as encryption keys are handled and managed by AWS. But is not what we're saying about "AWS managed key", so it will not automatic rotation every year.

<https://catalog.us-east-1.prod.workshops.aws/workshops/aad9ff1e-b607-45bc-893f-121ea5224f24/en-US/s3/serverside/sses3>

"AWS managed keys" are "KMS keys" in your account. And will (required) automatic rotation every year.

<https://docs.aws.amazon.com/kms/latest/developerguide/concepts.html#key-mgmt>

upvoted 2 times

FlyingHawk 4 months, 2 weeks ago

You are correct. SSE-S3 is not related to AWS KMS or AWS managed keys.

upvoted 1 times

EdricHoang 10 months, 2 weeks ago

there are several pages (not the official aws page) said its 365 days. But, the official page does not mention about the rotation period is 365 days

upvoted 1 times

Linuslin 1 year ago

SSE-KMS is similar to SSE-S3 but comes with some additional benefits over SSE-S3. And SSE-KMS is "AWS managed key." So it will (required) automatic rotation every year.

<https://catalog.us-east-1.prod.workshops.aws/workshops/aad9ff1e-b607-45bc-893f-121ea5224f24/en-US/s3/serverside/ssekms>

Difference between AWS S3 Bucket Encryption SSE-C , SSE-S3, SSE-KMS.

<https://awstip.com/5-minutes-to-aws-s3-bucket-encryption-sse-c-sse-s3-sse-kms-e2fb07b05cb3>

upvoted 1 times

TwinSpark 1 year ago

Selected Answer: B

I will go for B.

A it's somehow wrong for coupl of reason:

1- Encryption must be specified before to transfer the data (even if from 1/23 it's automatically for every bucket, so actually make no sense to specify it)

2- SSE-S3 keys are regularly rotated but aws do not specify when

(<https://docs.aws.amazon.com/AmazonS3/latest/userguide/serv-side-encryption.html>)

IMO if need to be compliance with rotation period better use Costumer managed key as stated from aws support in 01/2024

https://repost.aws/questions/QUES_1VN01TU-eRSO3LXergA/s3-managed-key-sse-s3-rotation-period

upvoted 2 times

bujuman 1 year ago

Selected Answer: A

Considering the statement "the LEAST operational overhead" we could go for option A due to the following AWS managed keys capabilities

https://docs.aws.amazon.com/kms/latest/developerguide/concepts.html#master_keys

upvoted 1 times

f04dc74 1 year, 1 month ago

Option A

upvoted 1 times

sandordini 1 year, 1 month ago

Selected Answer: A

From May 2022 the scheduled rotation is 1 year (SSE-S3)

upvoted 2 times

3b196fc 1 year, 1 month ago

A is wrong because you need to set the encryp options before send the data to S3.

upvoted 1 times

camps 1 year, 2 months ago

It's B.

upvoted 1 times

TruthWS 1 year, 2 months ago

A is correct because SSE-S3 help decrease the management

upvoted 2 times

Yushib 1 year, 2 months ago

Selected Answer: B

B is the right one

upvoted 2 times

haci 1 year, 2 months ago

Same with Question #202, I'll go with B but not sure
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #826

Topic 1

A company is migrating applications from an on-premises Microsoft Active Directory that the company manages to AWS. The company deploys the applications in multiple AWS accounts. The company uses AWS Organizations to manage the accounts centrally.

The company's security team needs a single sign-on solution across all the company's AWS accounts. The company must continue to manage users and groups that are in the on-premises Active Directory.

Which solution will meet these requirements?

- A. Create an Enterprise Edition Active Directory in AWS Directory Service for Microsoft Active Directory. Configure the Active Directory to be the identity source for AWS IAM Identity Center.
- B. Enable AWS IAM Identity Center. Configure a two-way forest trust relationship to connect the company's self-managed Active Directory with IAM Identity Center by using AWS Directory Service for Microsoft Active Directory. **Most Voted**
- C. Use AWS Directory Service and create a two-way trust relationship with the company's self-managed Active Directory.
- D. Deploy an identity provider (IdP) on Amazon EC2. Link the IdP as an identity source within AWS IAM Identity Center.

Correct Answer: B

Community vote distribution

B (100%)

Comments

LeonSauveterre 5 months, 1 week ago

Selected Answer: B

AWS IAM Identity Center (formerly AWS SSO): Provides SSO across multiple AWS accounts in an organization. By configuring a 2-way forest trust relationship between the on-premises AD and MAD, IAM Identity Center can integrate seamlessly with the existing AD to manage authentication and authorization, allowing the company to retain its existing on-premises AD as the primary identity source while extending to AWS.

A - This creates a separate directory in AWS, requiring migration of users and groups, which fails to continue using the on-premises AD.

C - IAM Identity Center is still needed for SSO (single sign-on) functionality.

D - OK but so much overhead.

upvoted 2 times

LuongTo 7 months, 1 week ago

why C is out?

upvoted 1 times

EdricHoang 11 months, 2 weeks ago

Selected Answer: B

"continue to manage users and groups that are in the on-premises Active Directory"

I go for B

upvoted 2 times

Scheldon 12 months ago

Selected Answer: B

AnswerB

AWS Directory Service lets you run Microsoft Active Directory (AD) as a managed service. AWS Directory Service for Microsoft Active Directory, also referred to as AWS Managed Microsoft AD, is powered by Windows Server 2019.

With AWS Managed Microsoft AD, you can run directory-aware workloads in the AWS Cloud, including Microsoft SharePoint and custom .NET and SQL Server-based applications. You can also configure a trust relationship between AWS Managed Microsoft AD in the AWS Cloud and your existing on-premises Microsoft Active Directory, providing users and groups with access to resources in either domain, using AWS IAM Identity Center.

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/directory_microsoft_ad.html

upvoted 2 times

Kaula 1 year, 2 months ago

Selected Answer: B

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/ms_ad_setup_trust.html

upvoted 4 times

haci 1 year, 2 months ago

Selected Answer: B

Same with Q-28

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #827

Topic 1


A company is planning to deploy its application on an Amazon Aurora PostgreSQL Serverless v2 cluster. The application will receive large amounts of traffic. The company wants to optimize the storage performance of the cluster as the load on the application increases.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure the cluster to use the Aurora Standard storage configuration.
- B. Configure the cluster storage type as Provisioned IOPS.
- C. Configure the cluster storage type as General Purpose.
- D. Configure the cluster to use the Aurora I/O-Optimized storage configuration. **Most Voted**

Correct Answer: D

Community vote distribution

**Comments**

Scheldon **Highly Voted** 12 months ago

Selected Answer: D

AnswerD

For Aurora we have 2 storage type/options:

A - Standard

B -- I/O-Optimized

hence answers B and C options are incorrect.

Because customer inform us that there will be a big amount of traffic for their application I would go with I/O-Optimized. Maybe we are giving more \$\$ per GB-month but we are not paying for I/O operations/request.

<https://aws.amazon.com/rds/aurora/pricing/>

In general question is not precise and it is hard to say which option will be more beneficial (cost effective)

upvoted 6 times

GOTJ **Most Recent** 3 months, 2 weeks ago

Selected Answer: A

In my opinion, when the company "wants to optimize the storage performance of the cluster AS the load on the application increases" they mean that they EXPECT that the load increases somewhere in the future. If so, what's the point of paying more for the I/O optimised storage configuration from the very beginning? You can switch to this more expensive configuration AS the load increased, saving money before this increment occurs

upvoted 1 times

mk168898 7 months, 3 weeks ago

Selected Answer: D

Aurora only have:

-> Standard

-> I/O-Optimized (need optimise storage thats why i chose this)

upvoted 2 times

Johnoppong101 9 months, 3 weeks ago

Selected Answer: D

Answer is D

upvoted 2 times

joseantoniopolo 1 year, 2 months ago

Selected Answer: D

Aurora I/O-Optimized – Improved price performance and predictability for I/O-intensive applications. You pay only for the usage and storage of your DB clusters, with no additional charges for read and write I/O operations.

upvoted 4 times

osmk 1 year, 2 months ago

Selected Answer: D

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.Overview.StorageReliability.html#aurora-storage-type>

upvoted 2 times

JCAWS 1 year, 2 months ago

Selected Answer: D

D is more suitable

upvoted 2 times

camps 1 year, 2 months ago

I would choose D

upvoted 2 times

TruthWS 1 year, 2 months ago

I think A is true answer

upvoted 2 times

xBUGx 1 year, 2 months ago

Selected Answer: D

<https://aws.amazon.com/about-aws/whats-new/2023/05/amazon-aurora-i-o-optimized/>

Aurora I/O-Optimized offers up to 40% cost savings for I/O-intensive applications where I/O charges exceed 25% of the total Aurora database spend.

upvoted 3 times

Kaula 1 year, 2 months ago

Selected Answer: C

Agree with haci

upvoted 2 times

haci 1 year, 2 months ago

Selected Answer: C

The traffic load is not defined well enough to decide which storage type to use.

General Purpose (SSD) storage suits many workloads, including small to medium-sized databases and it is cost-effective.

Provisioned IOPS (PIOPS) storage is the highest-performing option available for RDS instances. With Provisioned IOPS storage, you can provision a specific amount of IOPS (input/output operations per second) based on your application's needs. But here we don't know the amount of requests.

So since the question is asking for cost-effective I'll go with C
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #828

Topic 1

A financial services company that runs on AWS has designed its security controls to meet industry standards. The industry standards include the National Institute of Standards and Technology (NIST) and the Payment Card Industry Data Security Standard (PCI DSS).

The company's third-party auditors need proof that the designed controls have been implemented and are functioning correctly. The company has hundreds of AWS accounts in a single organization in AWS Organizations. The company needs to monitor the current state of the controls across accounts.

Which solution will meet these requirements?

- A. Designate one account as the Amazon Inspector delegated administrator account from the Organizations management account. Integrate Inspector with Organizations to discover and scan resources across all AWS accounts. Enable Inspector industry standards for NIST and PCI DSS.
- B. Designate one account as the Amazon GuardDuty delegated administrator account from the Organizations management account. In the designated GuardDuty administrator account, enable GuardDuty to protect all member accounts. Enable GuardDuty industry standards for NIST and PCI DSS.
- C. Configure an AWS CloudTrail organization trail in the Organizations management account. Designate one account as the compliance account. Enable CloudTrail security standards for NIST and PCI DSS in the compliance account.
- D. Designate one account as the AWS Security Hub delegated administrator account from the Organizations management account. In the designated Security Hub administrator account, enable Security Hub for all member accounts. Enable Security Hub standards for NIST and PCI DSS. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: D

Security Hub: assess your AWS environment against security industry standards and best practices.

Upvoted 7 times

upvoted / times

mk168898 Most Recent 7 months, 3 weeks ago

Selected Answer: D

security industry standards => security hub
upvoted 2 times

Johnoppong101 9 months, 3 weeks ago

Selected Answer: D

NIST, PCI DSS Compliance + AWS accounts -> Security Hub
upvoted 2 times

Kaula 1 year, 2 months ago

Selected Answer: D

<https://docs.aws.amazon.com/securityhub/latest/userguide/what-is-securityhub.html>
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #829

Topic 1

A company uses an Amazon S3 bucket as its data lake storage platform. The S3 bucket contains a massive amount of data that is accessed randomly by multiple teams and hundreds of applications. The company wants to reduce the S3 storage costs and provide immediate availability for frequently accessed objects.

What is the MOST operationally efficient solution that meets these requirements?

- A. Create an S3 Lifecycle rule to transition objects to the S3 Intelligent-Tiering storage class. **Most Voted**
- B. Store objects in Amazon S3 Glacier. Use S3 Select to provide applications with access to the data.
- C. Use data from S3 storage class analysis to create S3 Lifecycle rules to automatically transition objects to the S3 Standard-Infrequent Access (S3 Standard-IA) storage class.
- D. Transition objects to the S3 Standard-Infrequent Access (S3 Standard-IA) storage class. Create an AWS Lambda function to transition objects to the S3 Standard storage class when they are accessed by an application.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**mk168898** 7 months, 3 weeks ago**Selected Answer: A**

selected A because only option that don't violate the "need immediate access"
upvoted 2 times

MatAlves 8 months, 3 weeks ago**Selected Answer: A**

"data that is accessed randomly" = S3 Intelligent-Tiering storage class.
upvoted 2 times

Scheldon 12 months ago**Selected Answer: A**

AnswerA

The Amazon S3 Intelligent-Tiering storage class automatically stores objects in three access tiers. One tier is optimized for frequent access, one lower-cost tier is optimized for infrequent access, and another very low-cost tier is optimized for rarely accessed data. For a low monthly object monitoring and automation charge, S3 Intelligent-Tiering monitors access patterns and automatically moves objects to the Infrequent Access tier when they haven't been accessed for 30 consecutive days. After 90 days of no access, the objects are moved to the Archive Instant Access tier without performance impact or operational overhead.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/intelligent-tiering-overview.html>

upvoted 3 times

Hkayne 1 year ago

Selected Answer: A

File accessed randomly by multiple teams = intelligent tiering

upvoted 2 times

Kaula 1 year, 2 months ago

Selected Answer: A

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/intelligent-tiering-managing.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #830

Topic 1

A company has 5 TB of datasets. The datasets consist of 1 million user profiles and 10 million connections. The user profiles have connections as many-to-many relationships. The company needs a performance efficient way to find mutual connections up to five levels.

Which solution will meet these requirements?

- A. Use an Amazon S3 bucket to store the datasets. Use Amazon Athena to perform SQL JOIN queries to find connections.
- B. Use Amazon Neptune to store the datasets with edges and vertices. Query the data to find connections. **Most Voted**
- C. Use an Amazon S3 bucket to store the datasets. Use Amazon QuickSight to visualize connections.
- D. Use Amazon RDS to store the datasets with multiple tables. Perform SQL JOIN queries to find connections.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Kaula **Highly Voted** 1 year, 2 months ago

Selected Answer: B

<https://docs.aws.amazon.com/neptune/latest/userguide/notebooks-visualization.html>

upvoted 5 times

mk168898 **Most Recent** 7 months, 3 weeks ago

Selected Answer: B

everytime i see some social network related qns i immediately look for amazon neptune

upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: B

Neptune: A Graph database stores nodes and relationships instead of tables or documents

upvoted 3 times

alawada 1 year, 2 months ago

Selected Answer: B

Neptune automatically scales storage and compute resources based on workload demands, ensuring optimal performance even as the dataset grows over time.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #831

Topic 1

A company needs a secure connection between its on-premises environment and AWS. This connection does not need high bandwidth and will handle a small amount of traffic. The connection should be set up quickly.

What is the MOST cost-effective method to establish this type of connection?

- A. Implement a client VPN.
- B. Implement AWS Direct Connect.
- C. Implement a bastion host on Amazon EC2.
- D. Implement an AWS Site-to-Site VPN connection. **Most Voted**

Correct Answer: D*Community vote distribution*

D (100%)

Comments**mk168898** 7 months, 3 weeks ago**Selected Answer: D**

Not A because for individual connection not company

Not B because overkill, usually for high bandwidth, but qns clearly stated no need for high bandwidth and handle small traffic

Not C because bastion host for remote access

D because ideal for small amt of traffic

upvoted 2 times

Scheldon 1 year ago**Selected Answer: D**

AnswerD

AWS site-to-site VPN is the best solution here.

https://docs.aws.amazon.com/vpn/latest/s2svpn/VPC_VPN.html

upvoted 4 times

sandordini 1 year, 1 month ago**Selected Answer: D**

You can enable access to your remote (on-prem) network from your VPC by creating an AWS Site-to-Site VPN (Site-to-Site VPN) connection and configuring routing to pass traffic through the connection.

upvoted 3 times

Kaula 1 year, 2 months ago

Selected Answer: D

https://docs.aws.amazon.com/vpn/latest/s2svpn/VPC_VPN.html

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #832

Topic 1

A company has an on-premises SFTP file transfer solution. The company is migrating to the AWS Cloud to scale the file transfer solution and to optimize costs by using Amazon S3. The company's employees will use their credentials for the on-premises Microsoft Active Directory (AD) to access the new solution. The company wants to keep the current authentication and file access mechanisms.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Configure an S3 File Gateway. Create SMB file shares on the file gateway that use the existing Active Directory to authenticate.
- B. Configure an Auto Scaling group with Amazon EC2 instances to run an SFTP solution. Configure the group to scale up at 60% CPU utilization.
- C. Create an AWS Transfer Family server with SFTP endpoints. Choose the AWS Directory Service option as the identity provider. Use AD Connector to connect the on-premises Active Directory. **Most Voted**
- D. Create an AWS Transfer Family SFTP endpoint. Configure the endpoint to use the AWS Directory Service option as the identity provider to connect to the existing Active Directory.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Scheldon** 1 year ago**Selected Answer: C**

Answer C

We need Transfer Family SFTP enabled server (with SFTP endpoint). Additionally AWS Directory Service with AD connector to reach on-premises AD for authentication and authorization.

<https://docs.aws.amazon.com/transfer/latest/userguide/getting-started.html>

<https://docs.aws.amazon.com/transfer/latest/userguide/create-server-sftp.html>

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/what_is.html

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/directory_ad_connector.html

upvoted 3 times

shawn2022 1 year ago

phoenix2023 1 year ago

what is the difference between C and D ???

upvoted 2 times

EdricHoang 10 months, 2 weeks ago

AD Connector - when the company already had the Active Directory, you need a connector

upvoted 6 times

sandordini 1 year, 1 month ago

Selected Answer: C

1. Create one or more AWS Managed Microsoft AD directories using the AWS Directory Service console.
2. Use the Transfer Family console to create a **server** that uses **AWS Managed Microsoft AD** as its identity provider.
3. Add access from one or more of your AWS Directory Service groups.
4. Although not required, we recommend that you test and verify user access.

upvoted 3 times

Kaula 1 year, 2 months ago

Selected Answer: C

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/directory_ad_connector.html

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #833

Topic 1

A company is designing an event-driven order processing system. Each order requires multiple validation steps after the order is created. An idempotent AWS Lambda function performs each validation step. Each validation step is independent from the other validation steps. Individual validation steps need only a subset of the order event information.

The company wants to ensure that each validation step Lambda function has access to only the information from the order event that the function requires. The components of the order processing system should be loosely coupled to accommodate future business changes.

Which solution will meet these requirements?

- A. Create an Amazon Simple Queue Service (Amazon SQS) queue for each validation step. Create a new Lambda function to transform the order data to the format that each validation step requires and to publish the messages to the appropriate SQS queues. Subscribe each validation step Lambda function to its corresponding SQS queue.
- B. Create an Amazon Simple Notification Service (Amazon SNS) topic. Subscribe the validation step Lambda functions to the SNS topic. Use message body filtering to send only the required data to each subscribed Lambda function.
- C. Create an Amazon EventBridge event bus. Create an event rule for each validation step. Configure the input transformer to send only the required data to each target validation step Lambda function. **Most Voted**
- D. Create an Amazon Simple Queue Service (Amazon SQS) queue. Create a new Lambda function to subscribe to the SQS queue and to transform the order data to the format that each validation step requires. Use the new Lambda function to perform synchronous invocations of the validation step Lambda functions in parallel on separate threads.

Correct Answer: C

Community vote distribution

C (94%)

D (6%)

Comments

LeonSauveterre 5 months, 1 week ago

Selected Answer: C

A - Requires a separate Lambda function to preprocess and publish transformed data to multiple queues. Unnecessary.
B - SNS filtering operates at the subscription level, not the event transformation level, making it harder to customize data subsets for specific Lambda functions.

C - Amazon EventBridge: event-driven systems. Input Transformer: ensures each receives only the required subset of data.
D - Synchronous invocation adds tighter coupling between Lambda functions and the validation steps.

upvoted 2 times

MatAlves 8 months, 3 weeks ago

SNS and Message Filtering

- With SNS, message filtering allows you to control which subscribers receive messages based on attributes. However, the entire message is sent to each subscribed Lambda function; only those that match the filter criteria are processed.

EventBridge and Input Transformation

- EventBridge enables you to define rules that transform or modify events before they reach their targets. This allows you to customize the event payload, ensuring each validation step receives only the relevant information.

"The company wants to ensure that each validation step Lambda function has access to only the information from the order event that the function requires."

Therefore, C is the answer.

upvoted 4 times

Johnopping101 9 months, 3 weeks ago

Selected Answer: C

Event-driven Architecture + Each validation step needs ONLY a subset of the order EVENT created. Best way to transform this order even is EB Transformer.

upvoted 2 times

Scheldon 1 year ago

Selected Answer: C

Answer C

I wasn't sure but looks like EB with Input Transformation will allow for sending data which were choosed per destination

<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-pipes-input-transformation.html>

upvoted 3 times

TwinSpark 1 year ago

Selected Answer: C

not B because SNS cannot make messages manipulation, the option "message body filtering" will make discard or forward the FULL message if there is a matching field:

<https://docs.aws.amazon.com/sns/latest/dg/sns-message-filtering.html>

C - eventbus instead can manipulate event:

<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-event-bus.html>

D - Works, but too much operation IMO

upvoted 4 times

88f8032 1 year, 1 month ago

Why can't it be B?

upvoted 3 times

BBR01 1 year, 1 month ago

Selected Answer: D

It is D. It is one order event, not "events from many sources"

The main lambda parse the info to pieces, then makes synchronous invocations of the validation step Lambda functions on separate threads, and wait them to complete.

upvoted 1 times

waldirlsantos 1 year, 1 month ago

Selected Answer: C

IMO, C

"An event bus is a router that receives events and delivers them to zero or more destinations, or targets. Event buses are well-suited for routing events from many sources to many targets, with optional transformation of events prior to delivery to a target."

upvoted 2 times

upvoted 2 times

TruthWS 1 year, 2 months ago

Option C

upvoted 2 times

Kaula 1 year, 2 months ago

Selected Answer: C

<https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-event-bus.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #834

Topic 1

A company is migrating a three-tier application to AWS. The application requires a MySQL database. In the past, the application users reported poor application performance when creating new entries. These performance issues were caused by users generating different real-time reports from the application during working hours.

Which solution will improve the performance of the application when it is moved to AWS?

- A. Import the data into an Amazon DynamoDB table with provisioned capacity. Refactor the application to use DynamoDB for reports.
- B. Create the database on a compute optimized Amazon EC2 instance. Ensure compute resources exceed the on-premises database.
- C. Create an Amazon Aurora MySQL Multi-AZ DB cluster with multiple read replicas. Configure the application to use the reader endpoint for reports. **Most Voted**
- D. Create an Amazon Aurora MySQL Multi-AZ DB cluster. Configure the application to use the backup instance of the cluster as an endpoint for the reports.

Correct Answer: C

Community vote distribution

C (100%)

Comments

xBUGx **Highly Voted** 1 year, 2 months ago

Selected Answer: C

real-time reports - > read replica
upvoted 6 times

KennethNg923 **Most Recent** 11 months, 3 weeks ago

Selected Answer: C

caused by users generating different real-time reports -> read replica
upvoted 1 times

Scheldon 1 year ago

Selected Answer: C

Answer C

Any replica of Amazon Aurora Db will be read-only replica (even backup one). Option C is better than D because we will use multiple replicas which when used will significantly allow to increase performance for creating reports. In the same time no write operation should be affected.

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/AuroraMySQL.Replication.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #835

Topic 1

A company is expanding a secure on-premises network to the AWS Cloud by using an AWS Direct Connect connection. The on-premises network has no direct internet access. An application that runs on the on-premises network needs to use an Amazon S3 bucket.

Which solution will meet these requirements MOST cost-effectively?

- A. Create a public virtual interface (VIF). Route the AWS traffic over the public VIF.
- B. Create a VPC and a NAT gateway. Route the AWS traffic from the on-premises network to the NAT gateway.
- C. Create a VPC and an Amazon S3 interface endpoint. Route the AWS traffic from the on-premises network to the S3 interface endpoint. **Most Voted**
- D. Create a VPC peering connection between the on-premises network and Direct Connect. Route the AWS traffic over the peering connection.

Correct Answer: C*Community vote distribution***Comments****0c21057** 1 month, 3 weeks ago**Selected Answer: A**

AWS Direct Connect links your internal network to an AWS Direct Connect location over a standard Ethernet fiber-optic cable. One end of the cable is connected to your router, and the other to an AWS Direct Connect router. With this connection, you can create virtual interfaces directly to public AWS services (for example, to Amazon S3) or to Amazon VPC, bypassing internet service providers in your network path. An AWS Direct Connect location provides access to AWS in the Region with which it is associated. You can use a single connection in a public Region or AWS GovCloud (US) to access public AWS services in all other public Regions.

upvoted 1 times

Saliigen 5 months ago**Selected Answer: A**

<https://aws.amazon.com/blogs/networking-and-content-delivery/optimizing-amazon-s3-data-transfers-over-direct-connect/>

upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: C

- A - Public VIFs route traffic over the public internet, but on-premises network has no direct internet access.
- B - NAT gateways provide internet access for resources within a VPC, but, again, on-premises network has no direct internet access.
- C - The traffic stays within AWS and the Direct Connect connection, ensuring security and compliance with the no-internet requirement. Also, this is cost-effective because it avoids the need for additional NAT gateways or public VIFs.
- D - Peering connections don't even support access to AWS public services like S3.

upvoted 3 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: A

public VIF allows on-premises networks to access AWS public services, like S3
it eliminates the need for internet access & interface endpoints while being cost-effective

upvoted 2 times

EllenLiu 5 months, 2 weeks ago

<https://aws.amazon.com/privatelink/pricing/>

upvoted 1 times

AMEJack 6 months, 2 weeks ago

Selected Answer: A

Option A is the correct answer:

- 1) I need to connect to public services only, thus public VIF will be enough.
- 2) Interface endpoints is charged per hour, not cost effective.

upvoted 2 times

Scheldon 1 year ago

Selected Answer: C

AnswerC

Amazon S3 interface endpoint seems to be the best and only option as we are forced to use Private IP addressation.

Interface endpoints for Amazon S3

Your network traffic remains on the AWS network.

Use private IP addresses from your VPC to access Amazon S3

Require endpoint-specific Amazon S3 DNS names

Allow access from on premises

Allow access from a VPC in another AWS Region by using VPC peering or AWS Transit Gateway

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/privatelink-interface-endpoints.html#types-of-vpc-endpoints-for-s3>

upvoted 4 times

Obdf3af 1 year ago

A. <https://repost.aws/knowledge-center/s3-bucket-access-direct-connect>

upvoted 1 times

elmyth 9 months ago

This article says "Use a private IP address over Direct Connect (with an interface VPC endpoint)" - C

upvoted 3 times

Obdf3af 1 year ago

A. public VIF is the way you can connect on-premise with S3 via DirectConnect

upvoted 1 times

waldirlsantos 1 year, 1 month ago

Selected Answer: C

B Need internet

A,D doesn't connect to the s3

IMO C is the solution for this question

True, C is the solution for this question.

upvoted 4 times

TruthWS 1 year, 2 months ago

Option C

upvoted 2 times

Kaula 1 year, 2 months ago

Selected Answer: C

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/privatelink-interface-endpoints.html>

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #836

Topic 1

A company serves its website by using an Auto Scaling group of Amazon EC2 instances in a single AWS Region. The website does not require a database.

The company is expanding, and the company's engineering team deploys the website to a second Region. The company wants to distribute traffic across both Regions to accommodate growth and for disaster recovery purposes. The solution should not serve traffic from a Region in which the website is unhealthy.

Which policy or resource should the company use to meet these requirements?

- A. An Amazon Route 53 simple routing policy
- B. An Amazon Route 53 multivalue answer routing policy **Most Voted**
- C. An Application Load Balancer in one Region with a target group that specifies the EC2 instance IDs from both Regions
- D. An Application Load Balancer in one Region with a target group that specifies the IP addresses of the EC2 instances from both Regions

Correct Answer: B

Community vote distribution

B (100%)

Comments

waldirisantos **Highly Voted** 1 year, 1 month ago

Selected Answer: B

53 with multivalue is the best option for this case

Multivalue answer routing lets you configure Amazon Route 53 to return multiple values, such as IP addresses for your web servers, in response to DNS queries. You can specify multiple values for almost any record, but multivalue answer routing also lets you check the health of each resource, so Route 53 returns only values for healthy resources. It's not a substitute for a load balancer, but the ability to return multiple health-checkable IP addresses is a way to use DNS to improve availability and load balancing.

upvoted 5 times

mk168898 **Most Recent** 7 months, 3 weeks ago

Selected Answer: B

C and D are wrong because need serve traffic across both region.

B seems correct

upvoted 2 times

MatAlves 8 months, 3 weeks ago

A - Doesn't provide health check

C and D - Only work within a single zone.

upvoted 1 times

Sergiuss95 1 year, 1 month ago

Selected Answer: B

Yes, is the option b.

upvoted 1 times

TruthWS 1 year, 2 months ago

Option B

upvoted 1 times

Kaula 1 year, 2 months ago

Selected Answer: B

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-multivalue.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #837

Topic 1

A company runs its applications on Amazon EC2 instances that are backed by Amazon Elastic Block Store (Amazon EBS). The EC2 instances run the most recent Amazon Linux release. The applications are experiencing availability issues when the company's employees store and retrieve files that are 25 GB or larger. The company needs a solution that does not require the company to transfer files between EC2 instances. The files must be available across many EC2 instances and across multiple Availability Zones.

Which solution will meet these requirements?

- A. Migrate all the files to an Amazon S3 bucket. Instruct the employees to access the files from the S3 bucket.
- B. Take a snapshot of the existing EBS volume. Mount the snapshot as an EBS volume across the EC2 instances. Instruct the employees to access the files from the EC2 instances.
- C. Mount an Amazon Elastic File System (Amazon EFS) file system across all the EC2 instances. Instruct the employees to access the files from the EC2 instances. **Most Voted**
- D. Create an Amazon Machine Image (AMI) from the EC2 instances. Configure new EC2 instances from the AMI that use an instance store volume. Instruct the employees to access the files from the EC2 instances.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**freedafeng** **Highly Voted** 10 months, 3 weeks ago

my question is, why not A?
I am fine with C
upvoted 6 times

JA2018 6 months, 1 week ago

For Option A (S3 bucket): While S3 is great for storing large files, accessing files directly from S3 would require the employees to switch to a different interface, making it inconvenient and not aligned with the requirement to access files from the EC2 instances.

upvoted 2 times

NUC **Highly Voted** 1 year, 2 months ago

xBUGX Highly Voted 1 year, 2 months ago

Selected Answer: C

cross many EC2 instances and across multiple Availability Zones = EFS
upvoted 6 times

mk168898 Most Recent 7 months, 3 weeks ago

Selected Answer: C

"files must be available across many EC2 instances" means need some sort of shared system
immediately i look for EFS
upvoted 3 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: C

cross many EC2 instances and across multiple Availability Zones = EFS
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #838

Topic 1

A company is running a highly sensitive application on Amazon EC2 backed by an Amazon RDS database. Compliance regulations mandate that all personally identifiable information (PII) be encrypted at rest.

Which solution should a solutions architect recommend to meet this requirement with the LEAST amount of changes to the infrastructure?

- A. Deploy AWS Certificate Manager to generate certificates. Use the certificates to encrypt the database volume.
- B. Deploy AWS CloudHSM, generate encryption keys, and use the keys to encrypt database volumes.
- C. Configure SSL encryption using AWS Key Management Service (AWS KMS) keys to encrypt database volumes.
- D. Configure Amazon Elastic Block Store (Amazon EBS) encryption and Amazon RDS encryption with AWS Key Management Service (AWS KMS) keys to encrypt instance and database volumes. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

boluwatito **Highly Voted** 1 year, 1 month ago

Selected Answer: D

Amazon RDS relies on Amazon EBS volumes for storage.
By configuring Amazon EBS encryption, the underlying storage volumes are encrypted.

upvoted 6 times

JA2018 **Most Recent** 6 months, 1 week ago

Selected Answer: D

- Option D requires the least infrastructure modification as it leverages existing features of EBS and RDS to enable encryption with KMS keys.

- This is the most straightforward way to meet the compliance requirement without significant changes to the existing setup.

upvoted 1 times

mk168898 7 months, 3 weeks ago

SSL /Certificate -> encryption in transit so A and C are wrong

SSL/Certificate => encrypt in transit, so A and C are wrong.
so i feel the answer is between B and D.

upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: D

Encryption should be KMS, SSL is for transit not at rest...
Even though the question never mentioned any EBS volumes whatsoever, I would still go for D....

upvoted 4 times

zinabu 1 year, 1 month ago

answer:C

upvoted 2 times

zinabu 1 year, 1 month ago

answer:C

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #839

Topic 1

A company runs an AWS Lambda function in private subnets in a VPC. The subnets have a default route to the internet through an Amazon EC2 NAT instance. The Lambda function processes input data and saves its output as an object to Amazon S3.

Intermittently, the Lambda function times out while trying to upload the object because of saturated traffic on the NAT instance's network. The company wants to access Amazon S3 without traversing the internet.

Which solution will meet these requirements?

- A. Replace the EC2 NAT instance with an AWS managed NAT gateway.
- B. Increase the size of the EC2 NAT instance in the VPC to a network optimized instance type.
- C. Provision a gateway endpoint for Amazon S3 in the VPC. Update the route tables of the subnets accordingly. **Most Voted**
- D. Provision a transit gateway. Place transit gateway attachments in the private subnets where the Lambda function is running.

Correct Answer: C*Community vote distribution*

C (92%)

A (8%)

Comments**mk168898** 7 months, 3 weeks ago

without internet => gateway endpoints
upvoted 3 times

Scheldon 1 year ago**Selected Answer: C**

Answer C

"The company wants to access Amazon S3 without traversing the internet." so we cannot use any NAT like in answer A & B. Transit Gateways is allowing reach Direct Connect or VPN connection from VPC. Hence C need to be a good answer
upvoted 4 times

boubie44 1 year, 1 month ago

why not D? i don't understand
upvoted 2 times

DanielWuTRT 11 months, 2 weeks ago

Complexity and cost are high and too complicated for scenarios where only S3 access is required.
upvoted 2 times

waldirisantos 1 year, 1 month ago

Selected Answer: C

The Key words are "Without traversig the internet". So, the awser is C.
https://docs.aws.amazon.com/pt_br/vpc/latest/privatelink/gateway-endpoints.html
upvoted 4 times

AlvinC2024 1 year, 2 months ago

Selected Answer: C

By provisioning a gateway endpoint for Amazon S3 in the VPC, you enable the Lambda function running in the private subnets to access S3 directly without needing to go through the NAT instance or traverse the internet. This solution helps alleviate the network congestion issue and reduces latency since the traffic between Lambda and S3 stays within the AWS network. Additionally, updating the route tables of the subnets to route S3 traffic through the gateway endpoint ensures that the Lambda function can seamlessly communicate with S3 without encountering timeouts caused by network saturation on the NAT instance.
upvoted 3 times

dds69 1 year, 2 months ago

Selected Answer: A

NAT gateways are highly available and can automatically scale up to meet increased traffic demands.
upvoted 1 times

sandordini 1 year, 1 month ago

And uses the internet... So it can be a good solution, but not here, as: Without traversing the internet
upvoted 4 times

hpmargathia 1 year, 2 months ago

A
<https://aws.amazon.com/about-aws/whats-new/2015/12/introducing-amazon-vpc-nat-gateway-a-managed-nat-service/>
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #840

Topic 1

A news company that has reporters all over the world is hosting its broadcast system on AWS. The reporters send live broadcasts to the broadcast system. The reporters use software on their phones to send live streams through the Real Time Messaging Protocol (RTMP).

A solutions architect must design a solution that gives the reporters the ability to send the highest quality streams. The solution must provide accelerated TCP connections back to the broadcast system.

What should the solutions architect use to meet these requirements?

- A. Amazon CloudFront
- B. AWS Global Accelerator **Most Voted**
- C. AWS Client VPN
- D. Amazon EC2 instances and AWS Elastic IP addresses

Correct Answer: B

Community vote distribution

B (94%)

A (6%)

Comments

Mikado211 **Highly Voted** 1 year, 2 months ago

HTTP(S) -> Cloudfront
Other TCP -> AWS Global Accelerator
upvoted 5 times

Mikado211 1 year, 2 months ago

So the answer is B :)
upvoted 3 times

mk168898 **Most Recent** 7 months, 3 weeks ago

I see TCP/UDP related connection => AWS Global accelerator
upvoted 2 times

elmyth 9 months ago

Selected Answer: A

Looks like the question is very old, but before the right answer was Cloudfront. Now AWS says that "All RTMP workloads should begin migrating to a standard CloudFront Web distribution and use one of several HTTP streaming protocols such as HTTP Live Streaming (HLS), Dynamic Adaptive Streaming over HTTP (DASH), Microsoft Smooth Streaming (MSS), or HTTP Dynamic Streaming (HDS)." <https://repost.aws/questions/QUoUZgHZh7SEWlnQUPIBmVNQ/announcement-rtmp-support-discontinuing-on-december-31-2020>

upvoted 1 times

Scheldon 1 year ago

Selected Answer: B

AnswerB

We can eliminate C and D, A is for Web apps hence B should be ok

additionally <https://docs.aws.amazon.com/global-accelerator/latest/dg/introduction-how-it-works.html>

upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: B

Upload, TCP > AWS Global Accelerator

upvoted 4 times

Drew3000 1 year, 2 months ago

Last time I made it to the last question, they added more questions 2 mins later.

upvoted 2 times

Awsbeginner87 1 year, 2 months ago

They added 40 more questions today☐

upvoted 4 times

chasingsummer 1 year, 2 months ago

Selected Answer: B

Can't believe I finally made it to the last question. Good luck to everyone!

upvoted 4 times

TruthWS 1 year, 2 months ago

OptionB

upvoted 2 times

Kaula 1 year, 2 months ago

Where are questions 841-848?

I am I missing something?

upvoted 2 times

Kaula 1 year, 2 months ago

Selected Answer: B

B makes sense not A since CloudFront is CDN

upvoted 2 times

dds69 1 year, 2 months ago

Selected Answer: B

Global accelerator provides the acceleration for TCP

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #841

Topic 1

A company uses Amazon EC2 instances and Amazon Elastic Block Store (Amazon EBS) to run its self-managed database. The company has 350 TB of data spread across all EBS volumes. The company takes daily EBS snapshots and keeps the snapshots for 1 month. The daily change rate is 5% of the EBS volumes.

Because of new regulations, the company needs to keep the monthly snapshots for 7 years. The company needs to change its backup strategy to comply with the new regulations and to ensure that data is available with minimal administrative effort.

Which solution will meet these requirements MOST cost-effectively?

- A. Keep the daily snapshot in the EBS snapshot standard tier for 1 month. Copy the monthly snapshot to Amazon S3 Glacier Deep Archive with a 7-year retention period.
- B. Continue with the current EBS snapshot policy. Add a new policy to move the monthly snapshot to Amazon EBS Snapshots Archive with a 7-year retention period. **Most Voted**
- C. Keep the daily snapshot in the EBS snapshot standard tier for 1 month. Keep the monthly snapshot in the standard tier for 7 years. Use incremental snapshots.
- D. Keep the daily snapshot in the EBS snapshot standard tier. Use EBS direct APIs to take snapshots of all the EBS volumes every month. Store the snapshots in an Amazon S3 bucket in the Infrequent Access tier for 7 years.

Correct Answer: B

Community vote distribution



Comments

Scheldon **Highly Voted** 1 year ago

Selected Answer: B

AnswerB

The problem is that we need to choose best solution which is most cost-effective and have minimal administrative effort. Glacier is the best choice for 1st look, but there is one problem with that solution. From what I know there is no easy way to copy from EBS to Glacier and additionally current strategy is to make incremental snapshots. To copy file from EBS to (s3) Glacier we would need to run linux to which we will mount EBS and we will need copy everything to S3 and then move to glacier deep archive. And what is more you will have only incremental snapshot. Hence every solution which will say copy/move to S3 is not minimal administrative effort. Not mentioning that you will not have full snapshot

<https://stackoverflow.com/questions/41466044/amazon-ebs-snapshots-to-glacier>

<https://repost.aws/questions/QU5aCOBAtDK6WIMUz6BH3VqHA/move-ebs-to-glacier>

upvoted 9 times

MatAlves 8 months, 1 week ago

I first thought the same, but NO: "Which solution will meet these requirements MOST cost-effectively?" There is no mention to "operational overhead" or "minimal changes", etc. The question is asking purely "what is cheaper". So, A is the answer.

"We charge you \$0.0125 per GB-month of stored data and \$0.03 per GB retrieved. "

"All Storage / Month \$0.0036 per GB

S3 Glacier Deep Archive *** - For long-term data archiving that is accessed once or twice in a year and can be restored within 12 hours"

upvoted 2 times

elmyth 7 months, 1 week ago

"to ensure that data is available with minimal administrative effort." - there are 2 conditions

upvoted 4 times

rondellidell Highly Voted 1 year, 1 month ago

Selected Answer: A

How much does EBS snapshots archive cost?

Pricing and billing. Archived snapshots are billed at a rate of \$0.0125 per GB-month. For example, if you archive a 100 GiB snapshot, you are billed \$1.25 (100 GiB * \$0.0125) per month.

What is the cost of Glacier?

Even though uploading data to Amazon S3 Glacier is free, there is a pricing method for upload requests, which is \$0.03 per 1,000 requests. Transferring data out of S3 Glacier to the same region is free; however, there is a cost for transferring data to a different region.

- \$0.0036 per GB / Month

upvoted 7 times

TEC65 Most Recent 1 month, 1 week ago

Selected Answer: B

snapshots archive designed for low cost snapshots

upvoted 1 times

0c21057 1 month, 3 weeks ago

Selected Answer: B

The correct answer is B: Continue with the current EBS snapshot policy. Add a new policy to move the monthly snapshot to Amazon EBS Snapshots Archive with a 7-year retention period.

upvoted 1 times

gigilla 2 months, 1 week ago

Selected Answer: A

A: Amazon S3 Glacier Deep Archive

upvoted 1 times

nadeerm 3 months ago

Selected Answer: B

Cost Efficiency : Amazon EBS Snapshots Archive is designed to provide a cost-effective option for long-term storage of EBS snapshots. By moving monthly snapshots to the archive tier, you will significantly reduce the cost

upvoted 1 times

tch 3 months, 1 week ago

Selected Answer: B

Amazon EBS Snapshots Archive is a storage tier that you can use for low-cost, long-term storage of your rarely-accessed snapshots that do not need frequent or fast retrieval.

Move to S3... more job, so A is not a good idea...

Only 5% daily change would consider cheaper than S3 for long run (MOST cost-effectively)

Only 5% daily charges.. would consider cheaper than S3 for long run.. (MOST cost effectively)

I don't think in the exam need to calculate cost per \$, how to remember?

Just add new policy to move the snapshot to Amazon EBS Snapshots Archive (minimal administrative effort)
upvoted 1 times

Dantecito 3 months, 2 weeks ago

Selected Answer: A

A or B but I will choose A because Is the MOST cost-effectively and companies will always choose to save money.
upvoted 1 times

Salilgen 5 months ago

Selected Answer: B

"EBS Snapshots are stored incrementally."

"The snapshots are automatically saved to Amazon Simple Storage Service (Amazon S3) for long-term retention"

"Save up to 75% in snapshot storage costs by using EBS Snapshots Archive for the long-term retention (over 90 days) of seldom-accessed snapshots."

Moreover, copy files (option A) requires more administrative effort than automatic policy.

<https://aws.amazon.com/ebs/snapshots/#:~:text=available%20and%20durable-,Amazon%20EBS%20Snapshots%20are%20a%20convenient%20way%20to%20back%20up,availability%20of%20your%20EBS%20Snapshots.>
upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: A

A - S3 Glacier Deep Archive is designed for long-term data archiving with a retention period of up to 7 years or more at an extremely low cost (the most cost-effective S3 storage class).

B - OK but EBS Snapshots Archive costs more.

C - Keep the monthly snapshot in the standard tier for 1 month, OK. For 7 years??? You'll pay so much money for that. The standard EBS snapshot tier is designed for frequent access.

D - Storing snapshots directly in Amazon S3 (even in the Infrequent Access tier) is not ideal for EBS snapshots because it doesn't leverage EBS snapshot mechanisms, which are designed for efficient storage and management of EBS backup data.

upvoted 2 times

ARV14 6 months, 1 week ago

Selected Answer: B

<https://repost.aws/knowledge-center/ebs-copy-snapshot-data-s3-create-volume>. =>I want to copy an Amazon Elastic Block Store (Amazon EBS) snapshot to my Amazon Simple Storage Service (Amazon S3) bucket. I also want to create Amazon EBS volumes from data that's stored in my S3 bucket.

Looks like B

Short description

When you create an EBS snapshot, it's automatically stored in an Amazon S3 bucket that AWS manages. You can copy snapshots within the same AWS Region, or from one Region to another. However, you can't copy snapshots to S3 buckets that you manage.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: B

If we consider both cost and administrative effort more closely:

EBS Snapshots Archive may be easier to manage but could incur higher costs over 7 years compared to S3 Glacier.

S3 Glacier, despite its complexity for initial transfers, could end up being more cost-effective in the long run, especially for large data volumes.

Ultimately, if minimizing costs is a primary concern and the organization can handle the initial complexity of transferring snapshots, using S3 Glacier (Option A) could still be worth considering.

So, while Option B is easier, if cost is a significant factor, Option A might be the better choice despite the additional administrative effort involved. It's a trade-off that depends on the organization's priorities regarding cost and operational simplicity.

upvoted 4 times

MatAlves 8 months, 1 week ago

"We charge you \$0.0125 per GB-month of stored data and \$0.03 per GB retrieved. "

<https://aws.amazon.com/blogs/aws/new-amazon-ebs-snapshots-archive/>

"All Storage / Month \$0.0036 per GB
S3 Glacier Deep Archive *** - For long-term data archiving that is accessed once or twice in a year and can be restored within 12 hours"

<https://aws.amazon.com/s3/pricing/>

2nd time reviewing this question and, yeah, "A" is the option that better meets the "cost" requirement.
upvoted 2 times

dhewa 9 months, 2 weeks ago

Selected Answer: A

Keyword here is cost. S3 Glacier deep archive is significantly cheaper than keeping snapshots in the EBS snapshot standard tier or even the EBS Snapshots Archive.
upvoted 4 times

Abdullah2004 9 months, 3 weeks ago

Selected Answer: A

A is correct answer
upvoted 3 times

Johnoppong101 9 months, 3 weeks ago

Selected Answer: B

Choose EBS Snapshot Archive when:

Data is associated with EBS volumes.
You need to maintain point-in-time copies of your EBS volumes.
You require faster restore times than S3 Glacier Archive.
You need to comply with regulations requiring EBS snapshot retention.
upvoted 3 times

Zahran23 10 months, 4 weeks ago

Selected Answer: A

Option (B) is incorrect due to the following:
Archiving is recommended for monthly, quarterly, or yearly snapshots. Archiving daily incremental snapshots of a single volume can lead to higher costs when compared to keeping them in the standard tier.

<https://docs.aws.amazon.com/ebs/latest/userguide/snapshot-archive.html>
upvoted 5 times

Cpso 6 months, 1 week ago

the company needs to keep the monthly snapshots for 7 years. not daily
upvoted 1 times

Lin878 11 months, 1 week ago

Selected Answer: A

I would like to vote "A". we have to focus on the cost as per question.
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #842

Topic 1

A company runs an application on several Amazon EC2 instances that store persistent data on an Amazon Elastic File System (Amazon EFS) file system. The company needs to replicate the data to another AWS Region by using an AWS managed service solution.

Which solution will meet these requirements MOST cost-effectively?

- A. Use the EFS-to-EFS backup solution to replicate the data to an EFS file system in another Region.
- B. Run a nightly script to copy data from the EFS file system to an Amazon S3 bucket. Enable S3 Cross-Region Replication on the S3 bucket.
- C. Create a VPC in another Region. Establish a cross-Region VPC peer. Run a nightly rsync to copy data from the original Region to the new Region.
- D. Use AWS Backup to create a backup plan with a rule that takes a daily backup and replicates it to another Region. Assign the EFS file system resource to the backup plan. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****sandordini** **Highly Voted** 1 year, 1 month ago**Selected Answer: D**

NOTA: EFS-to-EFS backup: You must deploy this solution in the same AWS Region as your source Amazon EFS Filesystem

Not B, C: Not a managed AWS Solution

D: AWS backup will do the job, and is managed service.

upvoted 8 times

EdricHoang **Highly Voted** 11 months, 3 weeks ago**Selected Answer: D**

A is oke but it says ""AWS managed service solution"

So I go for D

upvoted 5 times

LeonSauveterre **Most Recent** 5 months, 1 week ago

Selected Answer: D

While EFS-to-EFS replication (Option A) is great for real-time replication, Option D (AWS Backup) is the more cost-effective choice for long-term retention and administrative simplicity.

Also, I copied this from somewhere:

EFS-to-EFS Replication might not technically be an "AWS Managed Service" in the same sense as services like AWS Backup or Amazon RDS, but it is a solution supported by AWS. The process is automated in the sense that AWS handles the replication between EFS file systems, but it does not provide the same level of managed service (like backup solutions) that abstracts away storage, retention, and backup policy management.

However, the EFS-to-EFS solution is often classified under the broader umbrella of AWS managed solutions, since it directly integrates with EFS, but AWS Backup may provide a more clear-cut "AWS managed" approach with better control over backup policies and retention.

upvoted 1 times

Cpso 6 months, 1 week ago

Selected Answer: D

A-D : managed solution

Cost $D < A$

upvoted 1 times

Madushanka 7 months, 2 weeks ago

Answer D: AWS Backup is a managed service that handles backup operations. If AWS Backup is not available in your region, you can consider using EFS-to-EFS backup

upvoted 2 times

mk168898 7 months, 3 weeks ago

AWS backup is the only managed service so D

upvoted 2 times

NSA_Poker 11 months, 3 weeks ago

Selected Answer: D

(A,B,C) are eliminated. They are NOT managed service solutions.

(D) is correct. 'AWS Backup offers a cost-effective, FULLY MANAGED, policy-based service that simplifies data protection at scale.'

<https://aws.amazon.com/getting-started/hands-on/amazon-efs-backup-and-restore-using-aws-backup/>

upvoted 3 times

Rahulmaddy 12 months ago

Selected Answer: A

A is cheaper and less complicated to implement compared to D

upvoted 3 times

Nm55569 1 year ago

Selected Answer: A

AWS backup is not replicating. EFS replication is and it's managed in that you configure it and then it does the replication - no further actions required. It's also cheapest since it's free, you just pay for data transfer and storage

<https://docs.aws.amazon.com/efs/latest/ug/efs-replication.html>

<https://aws.amazon.com/blogs/aws/new-replication-for-amazon-elastic-file-system-efs/>

upvoted 4 times

Scheldon 1 year ago

Selected Answer: D

AnswerD

I think the most cost effective would solution presented in C, but hence in question it's clearly wrote that we should use AWS Managed Services Solution, hence I think we have no other choice then to choose AWS backup (Option D)

upvoted 2 times

Scheldon 1 year ago

Screenoon 1 year ago

<https://docs.aws.amazon.com/managedservices/latest/userguide/features.html>

upvoted 2 times

expouser 1 year ago

A: Replication is available in all AWS Regions in which EFS is available. To use replication in a Region that is disabled by default, you must first opt in to the Region. For more information, see Managing AWS Regions in the AWS General Reference Reference Guide. If you later opt out of a Region, Amazon EFS pauses all replication activities for the Region. To resume replication activities for the Region, you need to again opt in to the AWS Region.

<https://docs.aws.amazon.com/efs/latest/ug/efs-replication.html>

upvoted 1 times

0bdf3af 1 year ago

A. you can replicate to another Region

upvoted 1 times

BatVanyo 1 year, 1 month ago

Selected Answer: D

To me "an AWS managed service solution" automatically translates to AWS Backup.

...Can't say if this is cheaper than EFS replication tho.

upvoted 3 times

xBUGx 1 year, 2 months ago

Selected Answer: A

To replicate data from an Amazon Elastic File System (EFS) file system to another AWS Region, the MOST cost-effective solution would be to use EFS Replication. Here's why:

EFS Replication:

EFS Replication allows you to natively create a copy of your file system in an AWS Region or Availability Zone (AZ) of your choice.

It automatically and transparently copies your data from the source file system to the destination, maintaining an RPO (Recovery Point Objective) of 15 minutes for most file systems.

This solution is specifically designed for replicating EFS data across Regions, ensuring data resilience and protection.

There are no additional costs for using replication failback, and you pay for the usual replication and file system changes as described in Amazon EFS pricing¹².

EFS Replication is available in all AWS Regions where EFS is available¹.

upvoted 3 times

boluwatito 1 year, 1 month ago

But it is not a managed service

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #843


Topic 1

An ecommerce company is migrating its on-premises workload to the AWS Cloud. The workload currently consists of a web application and a backend Microsoft SQL database for storage.

The company expects a high volume of customers during a promotional event. The new infrastructure in the AWS Cloud must be highly available and scalable.

Which solution will meet these requirements with the LEAST administrative overhead?

- A. Migrate the web application to two Amazon EC2 instances across two Availability Zones behind an Application Load Balancer. Migrate the database to Amazon RDS for Microsoft SQL Server with read replicas in both Availability Zones.
- B. Migrate the web application to an Amazon EC2 instance that runs in an Auto Scaling group across two Availability Zones behind an Application Load Balancer. Migrate the database to two EC2 instances across separate AWS Regions with database replication.
- C. Migrate the web application to Amazon EC2 instances that run in an Auto Scaling group across two Availability Zones behind an Application Load Balancer. Migrate the database to Amazon RDS with Multi-AZ deployment. **Most Voted**
- D. Migrate the web application to three Amazon EC2 instances across three Availability Zones behind an Application Load Balancer. Migrate the database to three EC2 instances across three Availability Zones.

Correct Answer: C*Community vote distribution***Comments****LeonSauveterre** 5 months, 1 week ago**Selected Answer: C**

Initially I chose A but then again, "highly available and scalable" should indicate multi az and auto scaling group, so C is the only valid option.

upvoted 1 times

Scheldon 1 year ago**Selected Answer: C**

AnswerC

I would de-couple as many elements as possible with appropriate redundancy (HA), hence Auto-Scale in 2 AZ for EC2, Database in Multi-AZ and ALB in front of EC2. It will allow to increase amount of servers in case of need and will prevent from service unavailability in case if something fail ;)

upvoted 3 times

zinabu 1 year, 1 month ago

Selected Answer: C

yes "C" but it was better if it says Amazon RDS for Microsoft SQL Multi-AZ. any ways

upvoted 3 times

KennethNg923 11 months, 3 weeks ago

Agree, Options Only C has auto scaling and RDS, and RDS support sql server, so it could only be C

upvoted 2 times

rondellidell 1 year, 1 month ago

Selected Answer: C

only c

upvoted 3 times

[Removed] 1 year, 1 month ago

Selected Answer: C

HA - option C

upvoted 2 times

Awsbeginner87 1 year, 2 months ago

Selected Answer: A

Option a

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #844

Topic 1

A company has an on-premises business application that generates hundreds of files each day. These files are stored on an SMB file share and require a low-latency connection to the application servers. A new company policy states all application-generated files must be copied to AWS. There is already a VPN connection to AWS.

The application development team does not have time to make the necessary code modifications to move the application to AWS.


Which service should a solutions architect recommend to allow the application to copy files to AWS?

- A. Amazon Elastic File System (Amazon EFS)
- B. Amazon FSx for Windows File Server
- C. AWS Snowball

D. AWS Storage Gateway **Most Voted**

Correct Answer: D

Community vote distribution

**Comments**

FlyingHawk 5 months ago

Selected Answer: D

Only B and D support smb share, so A and C are out, the requirement is to copy the files at smb shares to AWS automatically. Storage Gateway automatically copies files to AWS (e.g., Amazon S3) in the background, meeting the requirement to copy all application-generated files to AWS.

upvoted 3 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: D

D is good to go. AWS Storage Gateway provides seamless integration between on-premises environments and AWS, and specifically, File Gateway can be used to expose an SMB interface that on-premises applications can use.

For B: While FSx does support SMB, it's optimized for AWS cloud workloads, not necessarily for on-premises systems. If you want an on-premises system to directly interact with an SMB share in AWS you would typically need a service that integrates

with an on-premises system to directly interact with an SMB share in AWS, you would typically need a service that integrates with the on-premises network seamlessly.

upvoted 3 times

ckhemani 6 months, 1 week ago

Selected Answer: D

Amazon FSx for Windows File Server is not used of transfer data. if it was Amazon Fsx Gateway, then it would have been been correct.

They are looking solution for copying not moving to which filesystem.

upvoted 3 times

dhewa 9 months, 2 weeks ago

Selected Answer: D

B is incorrect in this case because the development team doesn't have time to modify the application.

upvoted 2 times

EdricHoang 10 months, 3 weeks ago

Selected Answer: B

Amazon FSx for Windows File Server also support SMB and hybrid solution. And, no modification is needed.

upvoted 3 times

Nm55569 1 year ago

Selected Answer: B

Why not "Amazon FSx for Windows File Server"?

upvoted 3 times

ike001 11 months, 3 weeks ago

because we require a hybrid solution here

upvoted 1 times

EdricHoang 10 months, 3 weeks ago

Amazon FSx for Windows File Server also support hybrid solution

upvoted 1 times

XXXXXINN 8 months ago

..I am wondering where it says FSx for Windows File Server also support hybrid solution?

upvoted 1 times

Scheldon 1 year ago

Selected Answer: D

AnswerD

StorageGateway

<https://docs.aws.amazon.com/storagegateway/>

upvoted 3 times

jck202020 1 year, 1 month ago

AWS Storage Gateway provides a set of hybrid cloud storage services that offer on-premises access to virtually unlimited cloud storage. The File Gateway configuration of AWS Storage Gateway supports the SMB protocol (and NFS), enabling on-premises applications to seamlessly store and retrieve files in Amazon S3 using existing file system protocols. It fits perfectly for applications that need to continue operating without modification, while also adhering to the new policy of copying files to AWS.

Given these descriptions, Option D (AWS Storage Gateway) is the recommended service. It allows for a smooth integration by maintaining the existing SMB file-sharing capabilities and connects seamlessly to AWS through the VPN, enabling daily file transfers without significant changes to application code or infrastructure.

upvoted 3 times

Awsbeginner87 1 year, 2 months ago

Selected Answer: D

AWS Storage Gateway service enables hybrid storage between on-premises environments and the AWS Cloud. It provides low-latency performance by caching frequently accessed data on premises, while storing data securely and durably in Amazon cloud storage services.

upvoted 2 times

Kaula 1 year, 2 months ago

Selected Answer: D

<https://docs.aws.amazon.com/storagegateway/>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #845

Topic 1

A company has 15 employees. The company stores employee start dates in an Amazon DynamoDB table. The company wants to send an email message to each employee on the day of the employee's work anniversary.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Create a script that scans the DynamoDB table and uses Amazon Simple Notification Service (Amazon SNS) to send email messages to employees when necessary. Use a cron job to run this script every day on an Amazon EC2 instance.
- B. Create a script that scans the DynamoDB table and uses Amazon Simple Queue Service (Amazon SQS) to send email messages to employees when necessary. Use a cron job to run this script every day on an Amazon EC2 instance.
- C. Create an AWS Lambda function that scans the DynamoDB table and uses Amazon Simple Notification Service (Amazon SNS) to send email messages to employees when necessary. Schedule this Lambda function to run every day. **Most Voted**
- D. Create an AWS Lambda function that scans the DynamoDB table and uses Amazon Simple Queue Service (Amazon SQS) to send email messages to employees when necessary. Schedule this Lambda function to run every day.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**KennethNg923** 11 months, 3 weeks ago**Selected Answer: C**

Operational efficiency, script need to be run every time, and send email need to use SNS, so it should be C
upvoted 3 times

Scheldon 1 year ago**Selected Answer: C**

AnswerC

Deploying full instance is an overkill. Lambda should be enough + SNS to sent email. And it should be quite cheap
upvoted 3 times

Mikado211 1 year, 1 month ago

Selected Answer: C

SNS for sending mails

Lambda to scan the database + send the message to the SNS topic.

Using a script on a EC2 will add maintenance on both the EC2 and the script + cronjobs are not reliable and can be hard to monitor properly.

SO answer C !

upvoted 3 times

Mikado211 1 year, 1 month ago

SNS for sending mails

Lambda to scan the database + send the message to the SNS topic.

Using a script on a EC2 will add maintenance on both the EC2 and the script + cronjobs are not reliable and can be hard to monitor properly.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #846

Topic 1

A company's application is running on Amazon EC2 instances within an Auto Scaling group behind an Elastic Load Balancing (ELB) load balancer. Based on the application's history, the company anticipates a spike in traffic during a holiday each year. A solutions architect must design a strategy to ensure that the Auto Scaling group proactively increases capacity to minimize any performance impact on application users.

Which solution will meet these requirements?

- A. Create an Amazon CloudWatch alarm to scale up the EC2 instances when CPU utilization exceeds 90%.
- B. Create a recurring scheduled action to scale up the Auto Scaling group before the expected period of peak demand.
Most Voted
- C. Increase the minimum and maximum number of EC2 instances in the Auto Scaling group during the peak demand period.
- D. Configure an Amazon Simple Notification Service (Amazon SNS) notification to send alerts when there are autoscaling:EC2_INSTANCE_LAUNCH events.

Correct Answer: B

Community vote distribution

B (84%)

Other (16%)

Comments

Cpso 6 months, 1 week ago

Selected Answer: C

objective is to optimize autoscaling but being fix number of ec2: "design a strategy to ensure that the Auto Scaling group proactively increases capacity to minimize any performance impact on application users."

upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: B

the company anticipates a spike in traffic during a holiday each year -> schedule action

upvoted 3 times

Scheldon 1 year ago

Selected Answer: B

AnswerB

<https://medium.com/@damadhav/aws-scaling-reactive-vs-proactive-vs-predictive-2701ad6d48c9>
upvoted 4 times

7ce90e0 1 year, 1 month ago

Selected Answer: B

it says proactively
upvoted 4 times

zinabu 1 year, 1 month ago

Selected Answer: B

it needs a scheduled action for the yearly holiday peak traffic
upvoted 2 times

Mikado211 1 year, 1 month ago

Selected Answer: B

Since we know when we will have a peak of activity. A scheduled scaling is a good idea.
upvoted 3 times

zinabu 1 year, 1 month ago

selected answer: B
it needs a scheduled action for the yearly holiday peak traffic
upvoted 2 times

Hkayne 1 year, 1 month ago

Selected Answer: A

The answer IS A
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #847

Topic 1

A company uses Amazon RDS for PostgreSQL databases for its data tier. The company must implement password rotation for the databases.

Which solution meets this requirement with the LEAST operational overhead?

- A. Store the password in AWS Secrets Manager. Enable automatic rotation on the secret. **Most Voted**
- B. Store the password in AWS Systems Manager Parameter Store. Enable automatic rotation on the parameter.
- C. Store the password in AWS Systems Manager Parameter Store. Write an AWS Lambda function that rotates the password.
- D. Store the password in AWS Key Management Service (AWS KMS). Enable automatic rotation on the AWS KMS key.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Scheldon 6 months, 1 week ago

Selected Answer: A

AnswerA

"In Secrets Manager, you can set up automatic rotation for your secrets."

<https://docs.aws.amazon.com/secretsmanager/latest/userguide/rotating-secrets.html>

upvoted 4 times

88f8032 7 months, 1 week ago

Selected Answer: A

A. AWS Secrets Manager

upvoted 2 times

jck202020 7 months, 4 weeks ago

Option A (Store the password in AWS Secrets Manager and enable automatic rotation on the secret) is the best solution. It meets the requirements with the least operational overhead by leveraging built-in features specifically designed for managing and rotating database credentials securely.

invoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #848

Topic 1

A company runs its application on Oracle Database Enterprise Edition. The company needs to migrate the application and the database to AWS. The company can use the Bring Your Own License (BYOL) model while migrating to AWS. The application uses third-party database features that require privileged access.

A solutions architect must design a solution for the database migration.

Which solution will meet these requirements MOST cost-effectively?

A. Migrate the database to Amazon RDS for Oracle by using native tools. Replace the third-party features with AWS Lambda.

B. Migrate the database to Amazon RDS Custom for Oracle by using native tools. Customize the new database settings to support the third-party features. **Most Voted**

C. Migrate the database to Amazon DynamoDB by using AWS Database Migration Service (AWS DMS). Customize the new database settings to support the third-party features.

D. Migrate the database to Amazon RDS for PostgreSQL by using AWS Database Migration Service (AWS DMS). Rewrite the application code to remove the dependency on third-party features.

Correct Answer: B

Community vote distribution

B (100%)

Comments

Scheldon 6 months ago

Selected Answer: B

AnswerB

BYOL + third-party features/applications = RDS Custom. Hence customer is using Oracle so we should use RDS customer for Oracle

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/rds-custom.html>

<https://aws.amazon.com/blogs/aws/amazon-rds-custom-for-oracle-new-control-capabilities-in-database-environment/>

<https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/working-with-custom-oracle.html>

upvoted 3 times

zinabu 7 months, 2 weeks ago

Selected Answer: B

Amazon RDS Custom for Oracle by using native tools= to support the third-party features.
upvoted 1 times

Hkayne 7 months, 3 weeks ago

Selected Answer: B

B IS suitable for this use case : the use of BYOL and the use of third party features with privileged access
upvoted 1 times

jcck202020 7 months, 4 weeks ago

Considering the requirements and the need to use Oracle Database features with privileged access and BYOL, Option B (Migrate the database to Amazon RDS Custom for Oracle by using native tools. Customize the new database settings to support the third-party features) is the most cost-effective and suitable solution. It allows for significant customization needed to accommodate specific third-party features while leveraging existing Oracle licenses.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #849

Topic 1

A large international university has deployed all of its compute services in the AWS Cloud. These services include Amazon EC2, Amazon RDS, and Amazon DynamoDB. The university currently relies on many custom scripts to back up its infrastructure. However, the university wants to centralize management and automate data backups as much as possible by using AWS native options.

Which solution will meet these requirements?

- A. Use third-party backup software with an AWS Storage Gateway tape gateway virtual tape library.
- B. Use AWS Backup to configure and monitor all backups for the services in use. **Most Voted**
- C. Use AWS Config to set lifecycle management to take snapshots of all data sources on a schedule.
- D. Use AWS Systems Manager State Manager to manage the configuration and monitoring of backup tasks.

Correct Answer: B*Community vote distribution*
B (100%)**Comments****Mikado211** **Highly Voted** 8 months, 1 week ago**Selected Answer: B**

Centralized management of backups == AWS Backup
upvoted 5 times

Scheldon **Most Recent** 6 months ago**Selected Answer: B**

AnswerB

AWS native tool that will support EC2, RDS and DynamoDB = AWS Backup

<https://docs.aws.amazon.com/aws-backup/latest/devguide/working-with-supported-services.html>
upvoted 3 times

Hkayne 7 months, 3 weeks ago**Selected Answer: B**

Automate backups = AWS Backup
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #850

Topic 1

A company wants to build a map of its IT infrastructure to identify and enforce policies on resources that pose security risks. The company's security team must be able to query data in the IT infrastructure map and quickly identify security risks.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon RDS to store the data. Use SQL to query the data to identify security risks.
- B. Use Amazon Neptune to store the data. Use SPARQL to query the data to identify security risks. **Most Voted**
- C. Use Amazon Redshift to store the data. Use SQL to query the data to identify security risks.
- D. Use Amazon DynamoDB to store the data. Use PartiQL to query the data to identify security risks.

Correct Answer: B*Community vote distribution***Comments****NSA_Poker** 11 months, 3 weeks ago**Selected Answer: B**

Visualize your AWS Infrastructure with Amazon Neptune and AWS Config
<https://aws.amazon.com/blogs/database/visualize-your-aws-infrastructure-with-amazon-neptune-and-aws-config/>

Using Amazon Neptune for Security Graphs
<https://aws.amazon.com/neptune/security-graphs-on-aws/#:~:text=Using%20Amazon%20Neptune%20for%20Security%20Graphs>
upvoted 2 times

Rahulmaddy 12 months ago**Selected Answer: D**

Option B is very tough to impliment for such a simple usecase.
upvoted 1 times

Scheldon 1 year ago**Selected Answer: B**

AnswerB

Neptune is the only Graph database from all options which seems to be most suitable.
I'm not sure about SPARQL but hence Neptune is the best option hence SPARQL should be to.
<https://docs.aws.amazon.com/neptune/latest/userguide/intro.html>

upvoted 3 times

jck202020 1 year, 1 month ago

Option B (Use Amazon Neptune to store the data. Use SPARQL to query the data) is the most suitable choice. Neptune is purpose-built for storing and querying graph data, making it a natural fit for representing and querying the complex relationships inherent in an IT infrastructure map. Additionally, SPARQL is a powerful and efficient query language for graph databases, facilitating quick identification of security risks.

upvoted 4 times

dds69 1 year, 2 months ago

Selected Answer: B

Using Amazon Neptune with SPARQL, a query language for graph databases, allows the security team to easily query the data in the IT infrastructure map to identify security risks. SPARQL is specifically designed for querying graph data and allows for complex queries to traverse relationships between resources efficiently.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #851

Topic 1


A large company wants to provide its globally located developers separate, limited size, managed PostgreSQL databases for development purposes. The databases will be low volume. The developers need the databases only when they are actively working.

Which solution will meet these requirements MOST cost-effectively?

- A. Give the developers the ability to launch separate Amazon Aurora instances. Set up a process to shut down Aurora instances at the end of the workday and to start Aurora instances at the beginning of the next workday.
- B. Develop an AWS Service Catalog product that enforces size restrictions for launching Amazon Aurora instances. Give the developers access to launch the product when they need a development database.
- C. Create an Amazon Aurora Serverless cluster. Develop an AWS Service Catalog product to launch databases in the cluster with the default capacity settings. Grant the developers access to the product. **Most Voted**
- D. Monitor AWS Trusted Advisor checks for idle Amazon RDS databases. Create a process to terminate identified idle RDS databases.

Correct Answer: C

Community vote distribution



Comments

NSA_Poker **Highly Voted** 11 months, 3 weeks ago

Selected Answer: C

(A,B,D) eliminated. Aurora instances & Amazon RDS use On-Demand or Reserved INSTANCES. These are more expensive than a serverless solution.

(C) is correct. Amazon Aurora Serverless automatically starts up, shuts down & scales capacity up or down based on your application's needs; you pay only for capacity consumed.

upvoted 9 times

kelmryan1 **Highly Voted** 1 year, 1 month ago

Yes but its asking for the most cost effective. B would cause frustration for developers if it was terminated unexpectedly. The answer should be C so developers can easily access when they are needed and auto scales based on demand

upvoted 5 times

upvoted 0 times

NSA_Poker 11 months, 3 weeks ago

You're correct for the wrong reasons. An instance being 'terminated unexpectedly' & it having 'easily access" are qualities related to high availability & not cost effectiveness.
(B) is wrong because it's more \$\$\$.

upvoted 2 times

LeonSauveterre Most Recent 5 months, 1 week ago

Selected Answer: C

A - Aurora = Overkill and expensive.

B - Just like A.

C - Aurora Serverless scales based on usage and can automatically scale down to zero when the database is not in use. This is ideal for a development environment where databases are only needed during active work hours.

D - Money is not an issue here I think. The real problem is - there's no guarantee that the database would be automatically available when the developers need it. The need for a manual process to monitor and terminate idle databases introduces operational overhead and the potential for human error.

upvoted 3 times

mk168898 7 months, 3 weeks ago

Selected Answer: C

when i see "managed PostgreSQL databases", i immediately look for serverless

upvoted 3 times

b3b5fdd 7 months, 4 weeks ago

Selected Answer: B

I think the answer correct is B as the database should be managed by the developer and Each developer needs a separate DB.

upvoted 1 times

Johnoppong101 9 months, 3 weeks ago

Selected Answer: B

Requirement: Each developer needs a separate DB. In a cluster, DB is shared among developers.

upvoted 1 times

Salilgen 5 months ago

"There's no limit to the number of databases you can have in an Aurora Serverless cluster"

<https://aws.amazon.com/blogs/database/best-practices-for-working-with-amazon-aurora-serverless/>

upvoted 1 times

Johnoppong101 9 months, 3 weeks ago

Guys, listen to the question, each developer -> a separate managed PostgreSQL DB.

this is the task. Then Most cost-effective comes. If you get a cheaper thing that does not fulfill the task, you are wrong.

Question: Can each developer get a separate DB for himself from the cluster?

upvoted 1 times

ike001 11 months, 3 weeks ago

Answer C: We know we require Service Catalog. This option provides workflow no how it works

upvoted 2 times

Scheldon 1 year ago

AnswerC

Using Aurora Serverless solution + AWS Service Catalog features seems to be a good idea

<https://aws.amazon.com/rds/aurora/serverless/>

<https://aws.amazon.com/servicecatalog/features/>

upvoted 2 times

f07ed8f 1 year ago

Selected Answer: B

Only option B can have limited size database
upvoted 1 times

NSA_Poker 11 months, 3 weeks ago

"when creating your Aurora database cluster, specify the desired range of database capacity or use the defaults, and connect your applications." This way we limit the capacity that we pay for.

<https://aws.amazon.com/rds/aurora/serverless/>
upvoted 5 times

[Removed] 1 year ago

Selected Answer: C

Aurora Serverless = inicia, encerra e escala a capacidade automaticamente de acordo com as necessidades.
AWS Service Catalog = catalogo de serviços que usuarios podem utilizar, dentro das configurações permitidas.

<https://aws.amazon.com/pt/rds/aurora/serverless/>
<https://aws.amazon.com/pt/servicecatalog/>
upvoted 3 times

Sergiuss95 1 year, 1 month ago

Selected Answer: C

I thin is c
upvoted 3 times

Hkayne 1 year, 1 month ago

Selected Answer: B

With AWS Service Catalog, you can meet your compliance requirements while making sure your customers can quickly deploy the cloud resources they need.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #852

Topic 1

A company is building a web application that serves a content management system. The content management system runs on Amazon EC2 instances behind an Application Load Balancer (ALB). The EC2 instances run in an Auto Scaling group across multiple Availability Zones. Users are constantly adding and updating files, blogs, and other website assets in the content management system.

A solutions architect must implement a solution in which all the EC2 instances share up-to-date website content with the least possible lag time.

Which solution meets these requirements?

A. Update the EC2 user data in the Auto Scaling group lifecycle policy to copy the website assets from the EC2 instance that was launched most recently. Configure the ALB to make changes to the website assets only in the newest EC2 instance.

B. Copy the website assets to an Amazon Elastic File System (Amazon EFS) file system. Configure each EC2 instance to mount the EFS file system locally. Configure the website hosting application to reference the website assets that are stored in the EFS file system. **Most Voted**

C. Copy the website assets to an Amazon S3 bucket. Ensure that each EC2 instance downloads the website assets from the S3 bucket to the attached Amazon Elastic Block Store (Amazon EBS) volume. Run the S3 sync command once each hour to keep files up to date.

D. Restore an Amazon Elastic Block Store (Amazon EBS) snapshot with the website assets. Attach the EBS snapshot as a secondary EBS volume when a new EC2 instance is launched. Configure the website hosting application to reference the website assets that are stored in the secondary EBS volume.

Correct Answer: B

Community vote distribution

B (100%)

Comments

mk168898 7 months, 3 weeks ago

Selected Answer: B

website content must be in sync with all EC2 instances, so use EFS
upvoted 1 times

ike001 11 months, 3 weeks ago

B is the answer
upvoted 1 times

Scheldon 1 year ago

Selected Answer: B

AnswerB

Looks like this is the only reasonable solution from all presented
upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: B

CMS is usually EFS
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #853

Topic 1

A company's web application consists of multiple Amazon EC2 instances that run behind an Application Load Balancer in a VPC. An Amazon RDS for MySQL DB instance contains the data. The company needs the ability to automatically detect and respond to suspicious or unexpected behavior in its AWS environment. The company already has added AWS WAF to its architecture.

What should a solutions architect do next to protect against threats?

- A. Use Amazon GuardDuty to perform threat detection. Configure Amazon EventBridge to filter for GuardDuty findings and to invoke an AWS Lambda function to adjust the AWS WAF rules. **Most Voted**
- B. Use AWS Firewall Manager to perform threat detection. Configure Amazon EventBridge to filter for Firewall Manager findings and to invoke an AWS Lambda function to adjust the AWS WAF web ACL.
- C. Use Amazon Inspector to perform threat detection and to update the AWS WAF rules. Create a VPC network ACL to limit access to the web application.
- D. Use Amazon Macie to perform threat detection and to update the AWS WAF rules. Create a VPC network ACL to limit access to the web application.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Awsbeginner87 **Highly Voted** 8 months, 1 week ago

Selected Answer: A

Malicious or suspicious activity - think of GuardDuty
upvoted 5 times

Scheldon **Most Recent** 6 months, 1 week ago

Selected Answer: A

AnswerA

Amazon GuardDuty is a threat detection service that continuously monitors for malicious activity and unauthorized behavior across your AWS environment.

aws guard duty environment
<https://aws.amazon.com/guardduty/features/>
upvoted 4 times

zinabu 7 months, 2 weeks ago

Selected Answer: A

Gard duty for automatic treat detection
upvoted 3 times

Hkayne 7 months, 3 weeks ago

Selected Answer: A

Threat detection means guardduty
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #854

Topic 1

A company is planning to run a group of Amazon EC2 instances that connect to an Amazon Aurora database. The company has built an AWS CloudFormation template to deploy the EC2 instances and the Aurora DB cluster. The company wants to allow the instances to authenticate to the database in a secure way. The company does not want to maintain static database credentials.

Which solution meets these requirements with the LEAST operational effort?

- A. Create a database user with a user name and password. Add parameters for the database user name and password to the CloudFormation template. Pass the parameters to the EC2 instances when the instances are launched.
- B. Create a database user with a user name and password. Store the user name and password in AWS Systems Manager Parameter Store. Configure the EC2 instances to retrieve the database credentials from Parameter Store.
- C. Configure the DB cluster to use IAM database authentication. Create a database user to use with IAM authentication. Associate a role with the EC2 instances to allow applications on the instances to access the database. **Most Voted**
- D. Configure the DB cluster to use IAM database authentication with an IAM user. Create a database user that has a name that matches the IAM user. Associate the IAM user with the EC2 instances to allow applications on the instances to access the database.

Correct Answer: C

Community vote distribution

C (100%)

Comments

Hkayne **Highly Voted** 7 months, 3 weeks ago

Selected Answer: C

Using IAM database authentication and associate a role with the EC2 instances is the least operational effort.

upvoted 6 times

Scheldon **Most Recent** 6 months, 1 week ago

Selected Answer: C

AnswerC

Customer would like to not manage password rotation from my understanding, hence A and B are not the best solution here. I don't think we can associate IAM user with EC2 instance, but we can associate IAM role.

In summary C

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_use_switch-role-ec2.html

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/iam-roles-for-amazon-ec2.html>

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/UsingWithRDS.IAMDBAuth.html>

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #855

Topic 1

A company wants to configure its Amazon CloudFront distribution to use SSL/TLS certificates. The company does not want to use the default domain name for the distribution. Instead, the company wants to use a different domain name for the distribution.

Which solution will deploy the certificate without incurring any additional costs?

- A. Request an Amazon issued private certificate from AWS Certificate Manager (ACM) in the us-east-1 Region.
- B. Request an Amazon issued private certificate from AWS Certificate Manager (ACM) in the us-west-1 Region.
- C. Request an Amazon issued public certificate from AWS Certificate Manager (ACM) in the us-east-1 Region. **Most Voted**
- D. Request an Amazon issued public certificate from AWS Certificate Manager (ACM) in the us-west-1 Region.

Correct Answer: C

Community vote distribution

C (100%)

Comments

LeonSauveterre 5 months, 1 week ago

Selected Answer: C

Yes, C. Just memorize that. For CloudFront distributions with custom domain names, ACM public certificates must be requested from us-east-1.

upvoted 3 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: C

Have to use east-1 region for ACM, and it should be public SSL/TLS for domain, so it should be C

upvoted 4 times

Scheldon 1 year ago

Selected Answer: C

AnswerC

Per AWS "Public SSL/TLS certificates provisioned through AWS Certificate Manager are free. You pay only for the AWS resources you create to run your application."

resources you create to run your application.
<https://aws.amazon.com/certificate-manager/pricing/?nc=sn&loc=3>

But hence AWS is recommending to use US east 1 I think I would go with C
Note

We recommend that you use ACM to provision, manage, and deploy SSL/TLS certificates on AWS managed resources. You must request an ACM certificate in the US East (N. Virginia) Region.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/cnames-and-https-procedures.html>

upvoted 3 times

sandordini 1 year, 1 month ago

Selected Answer: C

browsers trust public certificates automatically by default > C or D

To use an ACM certificate with Amazon CloudFront, you must request or import the certificate in the US East (N. Virginia) region [Nowhere is it stated why is this though...] > C

upvoted 2 times

BatVanyo 1 year, 1 month ago

Selected Answer: C

The certificate has to be public. The certificate has to be issued in us-east-1:

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/cnames-and-https-requirements.html>

"To use an ACM certificate with CloudFront, make sure you request (or import) the certificate in the US East (N. Virginia) Region (us-east-1)."

upvoted 3 times

rondellidell 1 year, 1 month ago

Selected Answer: C

<https://aws.amazon.com/certificate-manager/pricing/>

AWS Certificate Manager Pricing

Public SSL/TLS certificates provisioned through AWS Certificate Manager are free. You pay only for the AWS resources you create to run your application.

If you manage AWS Private Certificate Authority (CA) through ACM, refer to the AWS Private CA Pricing page for more details and examples.

upvoted 2 times

boluwatito 1 year, 1 month ago

Selected Answer: C

It is c

upvoted 3 times

boluwatito 1 year, 1 month ago

Should be c, it is a public certificate

upvoted 2 times

JackyCCK 1 year, 2 months ago

CloudFront should have a private cert and browser use public cert aiming to achieve non-repudiation. Ans should be A

upvoted 1 times

cloudee 1 year, 2 months ago

Selected Answer: C

This should be C. Private CA is not free

upvoted 2 times

Awsbeginner87 1 year, 2 months ago

Why not D. evrn option D is public CA

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #856

Topic 1

A company creates operations data and stores the data in an Amazon S3 bucket. For the company's annual audit, an external consultant needs to access an annual report that is stored in the S3 bucket. The external consultant needs to access the report for 7 days.

The company must implement a solution to allow the external consultant access to only the report.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Create a new S3 bucket that is configured to host a public static website. Migrate the operations data to the new S3 bucket. Share the S3 website URL with the external consultant.
- B. Enable public access to the S3 bucket for 7 days. Remove access to the S3 bucket when the external consultant completes the audit.
- C. Create a new IAM user that has access to the report in the S3 bucket. Provide the access keys to the external consultant. Revoke the access keys after 7 days.
- D. Generate a presigned URL that has the required access to the location of the report on the S3 bucket. Share the presigned URL with the external consultant. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Hkayne **Highly Voted** 7 months, 3 weeks ago

Selected Answer: D

Pre-signed URLs are used to provide short-term access to a private object in your S3 bucket. They work by appending an AWS Access Key, expiration time, and Sigv4 signature as query parameters to the S3 object.

upvoted 6 times

LeonSauveterre **Most Recent** 5 months, 1 week ago

Selected Answer: D

A - Yeah, too much trouble.

B - Even if you remove access after the 7 days, the bucket would be publicly accessible during the time frame, which could lead to accidental data exposure.

C - A little better than A, but you gotta create a new IAM user, provide access keys, and manually revoke them after 7 days. I'm not even sure I can remember that a week later...

D - No need to manage IAM users or keys. Also, everything you need to do can be done within 10 minutes.

upvoted 2 times

Scheldon 6 months, 1 week ago

Selected Answer: D

AnswerD

When using CLI to create presigned URL we can setup 7days (max) for URL expiration from the time of creation.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/ShareObjectPreSignedURL.html>

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #857

Topic 1

A company plans to run a high performance computing (HPC) workload on Amazon EC2 Instances. The workload requires low-latency network performance and high network throughput with tightly coupled node-to-node communication.

Which solution will meet these requirements?

- A. Configure the EC2 instances to be part of a cluster placement group. **Most Voted**
- B. Launch the EC2 instances with Dedicated Instance tenancy.
- C. Launch the EC2 instances as Spot Instances.
- D. Configure an On-Demand Capacity Reservation when the EC2 instances are launched.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Awsbeginner87 **Highly Voted** 1 year, 2 months ago

Selected Answer: A

Tightly coupled, low-latency,hpc - cluster placement group
upvoted 7 times

victor78 **Most Recent** 11 months, 3 weeks ago

I THINK IT'S C,D
upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: A

tightly coupled node-to-node communication -> placement group
upvoted 2 times

Scheldon 1 year ago

Selected Answer: A

AnswerA

I was thinking it will be D but after some research I think it will be A

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>
upvoted 3 times

Linuslin 1 year ago

Selected Answer: A

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/placement-groups.html>
upvoted 1 times

zinabu 1 year, 1 month ago

Selected Answer: A

he key word here is "tightly coupled node-to-node communication" which means we need to Configure the EC2 instances in a cluster placement group.
upvoted 1 times

Tanidanindo 1 year, 2 months ago

Selected Answer: A

all points to a cluster placement group
upvoted 3 times

AlvinC2024 1 year, 2 months ago

Selected Answer: A

The answer should be A.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #858

Topic 1

A company has primary and secondary data centers that are 500 miles (804.7 km) apart and interconnected with high-speed fiber-optic cable. The company needs a highly available and secure network connection between its data centers and a VPC on AWS for a mission-critical workload. A solutions architect must choose a connection solution that provides maximum resiliency.

Which solution meets these requirements?

- A. Two AWS Direct Connect connections from the primary data center terminating at two Direct Connect locations on two separate devices
- B. A single AWS Direct Connect connection from each of the primary and secondary data centers terminating at one Direct Connect location on the same device
- C. Two AWS Direct Connect connections from each of the primary and secondary data centers terminating at two Direct Connect locations on two separate devices **Most Voted**
- D. A single AWS Direct Connect connection from each of the primary and secondary data centers terminating at one Direct Connect location on two separate devices

Correct Answer: C

Community vote distribution

C (100%)

Comments

Hkayne **Highly Voted** 6 months, 2 weeks ago

Selected Answer: C

The most resilient is C
upvoted 5 times

Scheldon **Most Recent** 6 months, 1 week ago

Selected Answer: C

AnswerC

Maximum resilient solution will be to have 2 independent connections to AWS from each location (maximum redundancy)
upvoted 3 times

Q14001 6 months, 2 weeks ago

SN4321 6 months, 3 weeks ago

Selected Answer: C

Redundant Connections: Having two Direct Connect connections from each data center provides redundancy in case one connection fails.

Diverse Direct Connect Locations: Terminating the connections at two different Direct Connect locations further eliminates the risk of a single point of failure due to issues at a specific location.

Separate Devices: Using separate devices at each Direct Connect location adds another layer of redundancy, preventing a single device failure from impacting connectivity.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #859


Topic 1

A company runs several Amazon RDS for Oracle On-Demand DB instances that have high utilization. The RDS DB instances run in member accounts that are in an organization in AWS Organizations.

The company's finance team has access to the organization's management account and member accounts. The finance team wants to find ways to optimize costs by using AWS Trusted Advisor.

Which combination of steps will meet these requirements? (Choose two.)

- A. Use the Trusted Advisor recommendations in the management account. **Most Voted**
- B. Use the Trusted Advisor recommendations in the member accounts where the RDS DB instances are running.
- C. Review the Trusted Advisor checks for Amazon RDS Reserved Instance Optimization. **Most Voted**
- D. Review the Trusted Advisor checks for Amazon RDS Idle DB Instances.
- E. Review the Trusted Advisor checks for compute optimization. Crosscheck the results by using AWS Compute Optimizer.

Correct Answer: AC*Community vote distribution***Comments****Scheldon** **Highly Voted** 1 year ago**Selected Answer: AC**

AnswerAC

We have On-Demand instances of RDS. If DB is used very often then it can occur that using Reserved Instance can bring some \$\$ savings.

It is mentioned in AWS docs :)

<https://docs.aws.amazon.com/awssupport/latest/user/cost-optimization-checks.html#amazon-ec2-reserved-instances-optimization>

upvoted 5 times

7dcef09 **Most Recent** 1 month, 3 weeks ago**Selected Answer: AC**

A is correct

You can set up the organizational view feature to create a report for all member accounts in your AWS organization. For more information, see Organizational view for AWS Trusted Advisor.

upvoted 1 times

7dcef09 1 month, 3 weeks ago

<https://docs.aws.amazon.com/awssupport/latest/user/organizational-view.html>

upvoted 1 times

FlyingHawk 4 months, 4 weeks ago

Selected Answer: BC

A is incorrect, as Trust advisor recommendations are account specific. Trusted Advisor in the management account provides insights only for that account, not for member accounts. Unless you have a centralized tool that aggregates recommendations, this option alone does not meet the requirement. C helps identify opportunities to save money by purchasing Reserved Instances for RDS if usage patterns justify it. This is one of the key cost optimization strategies for highly utilized RDS instances.

upvoted 1 times

FlyingHawk 4 months, 4 weeks ago

"You can use the Trusted Advisor Recommendations page of the Trusted Advisor console to review check results for your AWS account and then follow the recommended steps to fix any issues. " Trust advisor recommendations are account specific.

upvoted 2 times

FlyingHawk 4 months, 4 weeks ago

<https://docs.aws.amazon.com/awssupport/latest/user/get-started-with-aws-trusted-advisor.html>

upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: AC

B is OK but less efficient than A.

D is extremely unlikely because "instances have high utilization".

E - The finance team is focused on RDS for Oracle On-Demand DB instances, not compute (EC2) instances, so this check won't be useful for optimizing RDS costs.

upvoted 1 times

MatAlves 8 months, 3 weeks ago

Selected Answer: AC

"On-Demand DB instances that have high utilization"

High Utilization = no point in checking for Idle instances

On-Demand = it makes sense to replace On-demand for Reserved instances.

upvoted 4 times

example_ 11 months ago

Selected Answer: CD

<https://docs.aws.amazon.com/awssupport/latest/user/cost-optimization-checks.html>

upvoted 1 times

victor78 11 months, 3 weeks ago

I THINK IT'S C,D

upvoted 2 times

example_ 11 months ago

<https://docs.aws.amazon.com/awssupport/latest/user/cost-optimization-checks.html>

upvoted 1 times

Tomrr 1 year ago

Selected Answer: AC

Instances are running On-Demand, they need to check to see if they can save money by switching the reserved instances from on-demand

An important part of using AWS involves balancing your Reserved Instance (RI) purchase against your On-Demand Instance usage. This check provides recommendations on which RIs will help reduce the costs incurred from using On-Demand Instances.

upvoted 3 times

bujuman 1 year ago

Selected Answer: AD

Cost Optimization: By providing actionable recommendations, Trusted Advisor assists you in identifying areas of overspending and underutilization, like idle RDS DB instances or underused EBS volumes, leading to significant cost savings.

upvoted 1 times

f07ed8f 1 year ago

Selected Answer: AD

Please check with question 308. It should run in management account and review the Trusted Advisor check for Amazon RDS Idle DB Instances.

upvoted 1 times

f07ed8f 1 year ago

I correct myself the answer should be AC as the DB can be switch to on-demand rather than reserved and it would save save the cost.

upvoted 2 times

rondellidell 1 year, 1 month ago

Selected Answer: AC

<https://docs.aws.amazon.com/awssupport/latest/user/cost-optimization-checks.html#amazon-rds-reserved-instance-optimization>

upvoted 4 times

Awsbeginner87 1 year, 2 months ago

Selected Answer: AE

Option A,E

upvoted 1 times

xBUGx 1 year, 2 months ago

i dont think AWS Compute Optimizer work with RDS

upvoted 1 times

xBUGx 1 year, 2 months ago

Selected Answer: AD

use Trusted advisor on management account

upvoted 1 times

AlvinC2024 1 year, 2 months ago

Selected Answer: AC

<https://docs.aws.amazon.com/awssupport/latest/user/organizational-view.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #860

Topic 1

A solutions architect is creating an application. The application will run on Amazon EC2 instances in private subnets across multiple Availability Zones in a VPC. The EC2 instances will frequently access large files that contain confidential information. These files are stored in Amazon S3 buckets for processing. The solutions architect must optimize the network architecture to minimize data transfer costs.

What should the solutions architect do to meet these requirements?

- A. Create a gateway endpoint for Amazon S3 in the VPC. In the route tables for the private subnets, add an entry for the gateway endpoint. **Most Voted**
- B. Create a single NAT gateway in a public subnet. In the route tables for the private subnets, add a default route that points to the NAT gateway.
- C. Create an AWS PrivateLink interface endpoint for Amazon S3 in the VPC. In the route tables for the private subnets, add an entry for the interface endpoint.
- D. Create one NAT gateway for each Availability Zone in public subnets. In each of the route tables for the private subnets, add a default route that points to the NAT gateway in the same Availability Zone.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**Scheldon** 6 months, 1 week ago**Selected Answer: A**

Answer A

I think only option A have any sense. It is cheap (no cost), it is secure (traffic is not going to public network).
<https://docs.aws.amazon.com/vpc/latest/privatelink/gateway-endpoints.html>
upvoted 3 times

Hkayne 7 months, 3 weeks ago**Selected Answer: A**

Aws gateway will have no cost because the traffic will stay on aws infrastructure.

the gateway will have no cost because the traffic will stay on the infrastructure.
upvoted 2 times

Tanidanindo 8 months ago

Selected Answer: A

Gateway endpoint will minimize data transfer costs
upvoted 2 times

Awsbeginner87 8 months, 1 week ago

Selected Answer: A

A- gateway endpoint for S3
upvoted 2 times

xBUGx 8 months, 1 week ago

Selected Answer: A

gateway endpoint for Amazon S3
upvoted 1 times

AlvinC2024 8 months, 1 week ago

Selected Answer: A

Gateway endpoint is free <https://digitalcloud.training/vpc-interface-endpoint-vs-gateway-endpoint-in-aws/>.
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #861

Topic 1


A company wants to relocate its on-premises MySQL database to AWS. The database accepts regular imports from a client-facing application, which causes a high volume of write operations. The company is concerned that the amount of traffic might be causing performance issues within the application.

How should a solutions architect design the architecture on AWS?

- A. Provision an Amazon RDS for MySQL DB instance with Provisioned IOPS SSD storage. Monitor write operation metrics by using Amazon CloudWatch. Adjust the provisioned IOPS if necessary. **Most Voted**
- B. Provision an Amazon RDS for MySQL DB instance with General Purpose SSD storage. Place an Amazon ElastiCache cluster in front of the DB instance. Configure the application to query ElastiCache instead.
- C. Provision an Amazon DocumentDB (with MongoDB compatibility) instance with a memory optimized instance type. Monitor Amazon CloudWatch for performance-related issues. Change the instance class if necessary.
- D. Provision an Amazon Elastic File System (Amazon EFS) file system in General Purpose performance mode. Monitor Amazon CloudWatch for IOPS bottlenecks. Change to Provisioned Throughput performance mode if necessary.

Correct Answer: A

Community vote distribution



Comments

Hkayne **Highly Voted** 1 year, 1 month ago

Selected Answer: A

A or B. Can't be B because there is high volume of write no need for ElastiCache
upvoted 6 times

GOTJ **Most Recent** 3 months, 3 weeks ago

Selected Answer: B

Even though option "A" makes sense, everybody discarding option "B" is focusing on improving the import process performance itself by using ElastiCache, which is basically correct. However, there is not a clear indicator that the "client-facing application" ONLY performs high volumes of write operations. In fact, this can be pointless (a "write only" database?).

In my opinion, the main concern of the company is how THE REST of the application processes might be affected while this

write-intensive operations run in the database. And those operations surely includes reads.

Using an ElastiCache in front of the RDB instance or creating a read replica to redirect reads to might help with performance of THE REST of application when write-intensive imports processes are executing, as the query result might be cached.

My vote goes to "B"
upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: A

high volume of write operation -> Provisioned IOPS SSD storage
upvoted 3 times

Scheldon 1 year ago

Selected Answer: A

Answer A

For sure we cannot choose generar purpose IOPS SSD hence I would choos provisioned one. additionally it is a good idea to monitor performance with CloudWatch and adjust setup(provisioned IOPS) if there will be a need.

upvoted 4 times

sandordini 1 year, 1 month ago

Selected Answer: B

The most effective strategy for coping with that limit is to supplement disk-based databases with in-memory caching (Elasticache for Redis, Write-through strategy) I'd go for B...

upvoted 3 times

JA2018 6 months, 1 week ago

why I would not choose #B? While ElastiCache can be used for caching read-heavy workloads, it's not the best choice for a database with high write operations as it is primarily designed for fast reads from a cache.

upvoted 1 times

Mr_Marcus 1 year ago

If it were changes to existing data, maybe. The scenario specifically says data imports. Going with "A".

upvoted 2 times

Tanidanindo 1 year, 2 months ago

Selected Answer: A

Amazon RDS for MySQL DB instance with Provisioned IOPS SSD storage
upvoted 4 times