

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #862

Topic 1

A company runs an application in the AWS Cloud that generates sensitive archival data files. The company wants to rearchitect the application's data storage. The company wants to encrypt the data files and to ensure that third parties do not have access to the data before the data is encrypted and sent to AWS. The company has already created an Amazon S3 bucket.

Which solution will meet these requirements?

- A. Configure the S3 bucket to use client-side encryption with an Amazon S3 managed encryption key. Configure the application to use the S3 bucket to store the archival files.
- B. Configure the S3 bucket to use server-side encryption with AWS KMS keys (SSE-KMS). Configure the application to use the S3 bucket to store the archival files.
- C. Configure the S3 bucket to use dual-layer server-side encryption with AWS KMS keys (SSE-KMS). Configure the application to use the S3 bucket to store the archival files.
- D. Configure the application to use client-side encryption with a key stored in AWS Key Management Service (AWS KMS). Configure the application to store the archival files in the S3 bucket. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

xBUGx **Highly Voted** 8 months ago

Selected Answer: D

"ensure that third parties do not have access to the data before the data is encrypted and sent to AWS"
upvoted 8 times

f07ed8f **Highly Voted** 6 months, 2 weeks ago

Selected Answer: D

"Amazon S3 managed encryption key" (SSE-S3) is a server-side encryption. Therefore it is not a client-side encryption. To encrypt the data before sending to S3, it has to be client-side encryption.
upvoted 6 times

Hkayne **Most Recent** 6 months, 2 weeks ago

Selected Answer: D

Must encrypt the data on client side before uploading it to S3
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #863

Topic 1

A company uses Amazon RDS with default backup settings for its database tier. The company needs to make a daily backup of the database to meet regulatory requirements. The company must retain the backups for 30 days.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Write an AWS Lambda function to create an RDS snapshot every day.
- B. Modify the RDS database to have a retention period of 30 days for automated backups. **Most Voted**
- C. Use AWS Systems Manager Maintenance Windows to modify the RDS backup retention period.
- D. Create a manual snapshot every day by using the AWS CLI. Modify the RDS backup retention period.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**Hkayne** **Highly Voted** 1 year, 1 month ago**Selected Answer: B**

By default, Amazon RDS creates and saves automated backups of your DB instance securely in Amazon S3 for a user-specified retention period. You can set the backup retention period from 1 to 35 days. The maximum retention period currently available for automated snapshots is 35 days. When automated backups are turned on for your DB Instance, Amazon RDS automatically performs a full, daily snapshot of your data and captures transaction logs.

upvoted 5 times

EdricHoang **Most Recent** 11 months, 3 weeks ago**Selected Answer: B**

Its B

"Amazon RDS performs a full daily backup of your data during a backup window that you define when you create the DB instance. You can configure a retention period of up to 35 days for the automated backup."

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #864

Topic 1


A company that runs its application on AWS uses an Amazon Aurora DB cluster as its database. During peak usage hours when multiple users access and read the data, the monitoring system shows degradation of database performance for the write queries. The company wants to increase the scalability of the application to meet peak usage demands.

Which solution will meet these requirements MOST cost-effectively?

- A. Create a second Aurora DB cluster. Configure a copy job to replicate the users' data to the new database. Update the application to use the second database to read the data.
- B. Create an Amazon DynamoDB Accelerator (DAX) cluster in front of the existing Aurora DB cluster. Update the application to use the DAX cluster for read-only queries. Write data directly to the Aurora DB cluster.
- C. Create an Aurora read replica in the existing Aurora DB cluster. Update the application to use the replica endpoint for read-only queries and to use the cluster endpoint for write queries. **Most Voted**
- D. Create an Amazon Redshift cluster. Copy the users' data to the Redshift cluster. Update the application to connect to the Redshift cluster and to perform read-only queries on the Redshift cluster.

Correct Answer: C

Community vote distribution



Comments

mk168898 7 months, 2 weeks ago

I was thinking between B and C. Saw that DAX is for Dynamo DB which is different from the question Aurora DB. So i picked C
upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: C

Read Replica for multiple users access and read the data
upvoted 4 times

ug56c 12 months ago

Answer C wording is that read queries are slowing down write queries => we need to optimize for read queries => we need to

Answer C, wording is that read queries are slowing down write queries -> we need to optimize for read queries -> we need to add read replicas.

<https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.Managing.Performance.html#Aurora.Managing.Performance.ReadScaling>

upvoted 4 times

sheilawu 1 year ago

Selected Answer: B

write queries=>High performance of read & Write Dynamo, I think DAX is a better solution.

upvoted 1 times

Rhydian25 11 months, 1 week ago

Dynamo is NoSQL...

upvoted 2 times

Scheldon 1 year ago

Selected Answer: C

AnswerC

A - We can create only RO replica of DB hence this is not possible.

D - Redshift is not for a DB from my understanding but more like analytics tools

B - Could be a thinkg but there is no point of Read-Only from DAX and write to DB cluster. beside it is hard to say if it would be cost effective as we are paying per hour not R/W requests as it is with replica.

Hence I would go with C

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #865

Topic 1

A company's near-real-time streaming application is running on AWS. As the data is ingested, a job runs on the data and takes 30 minutes to complete. The workload frequently experiences high latency due to large amounts of incoming data. A solutions architect needs to design a scalable and serverless solution to enhance performance.

Which combination of steps should the solutions architect take? (Choose two.)

- A. Use Amazon Kinesis Data Firehose to ingest the data. **Most Voted**
- B. Use AWS Lambda with AWS Step Functions to process the data.
- C. Use AWS Database Migration Service (AWS DMS) to ingest the data.
- D. Use Amazon EC2 instances in an Auto Scaling group to process the data.
- E. Use AWS Fargate with Amazon Elastic Container Service (Amazon ECS) to process the data. **Most Voted**

Correct Answer: AE

Community vote distribution

AE (100%)

Comments

xBUGx **Highly Voted** 1 year, 2 months ago

Selected Answer: AE

Lambda maxed to 15mins
upvoted 5 times

JoeTromundo 8 months, 2 weeks ago

But through Step Functions it would be possible to divide a long task, lasting more than 15 minutes, into multiple Lambda invocations with a maximum duration of 15 minutes.
upvoted 1 times

JA2018 6 months, 1 week ago

but that will introduce additional overhead.
upvoted 1 times

mk168898 **Most Recent** 7 months, 2 weeks ago

mr10000 7 months, 2 weeks ago

Eliminate B because lambda max handle 15minutes.

upvoted 2 times

Scheldon 1 year ago

Selected Answer: AE

AnswerAE

Amazon Kinesis data firehose for data ingesting, and hence output cannot go to EC2, hence Fargate with ECS.

upvoted 3 times

Hkayne 1 year, 1 month ago

Selected Answer: AE

A is correct for ingesting data.

B or E both choises are serverless but the difference is the lambda maximum execution time is 15 minutes. So the right option is E.

A and E

upvoted 3 times

Awsbeginner87 1 year, 2 months ago

Selected Answer: AE

A-ingesting real-time data

E- serverless option ECS+fargate

upvoted 4 times

AlvinC2024 1 year, 2 months ago

Selected Answer: AE

The maximum run time for lambda is 15 mins.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #866

Topic 1

A company runs a web application on multiple Amazon EC2 instances in a VPC. The application needs to write sensitive data to an Amazon S3 bucket. The data cannot be sent over the public internet.

Which solution will meet these requirements?

- A. Create a gateway VPC endpoint for Amazon S3. Create a route in the VPC route table to the endpoint. **Most Voted**
- B. Create an internal Network Load Balancer that has the S3 bucket as the target.
- C. Deploy the S3 bucket inside the VPC. Create a route in the VPC route table to the bucket.
- D. Create an AWS Direct Connect connection between the VPC and an S3 regional endpoint.

Correct Answer: A

Community vote distribution

A (100%)

Comments

mk168898 7 months, 2 weeks ago

not allowed to access internet, and need to access S3 => gateway VPC endpoint
upvoted 2 times

Scheldon 1 year ago

Selected Answer: A

Answer A
upvoted 2 times

sandordini 1 year, 1 month ago

Selected Answer: A

no internet, S3 > gateway VPC endpoint
upvoted 3 times

Tanidanindo 1 year, 2 months ago

Selected Answer: A

VPC endpoint

11 responses
upvoted 2 times

Mikado211 1 year, 2 months ago

Selected Answer: A

"data cannot be sent over the public internet." == VPC Endpoint
upvoted 2 times

Awsbeginner87 1 year, 2 months ago

Selected Answer: A

Option A
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #867

Topic 1


A company runs its production workload on Amazon EC2 instances with Amazon Elastic Block Store (Amazon EBS) volumes. A solutions architect needs to analyze the current EBS volume cost and to recommend optimizations. The recommendations need to include estimated monthly saving opportunities.

Which solution will meet these requirements?

- A. Use Amazon Inspector reporting to generate EBS volume recommendations for optimization.
- B. Use AWS Systems Manager reporting to determine EBS volume recommendations for optimization.
- C. Use Amazon CloudWatch metrics reporting to determine EBS volume recommendations for optimization.
- D. Use AWS Compute Optimizer to generate EBS volume recommendations for optimization. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

KennethNg923 11 months, 3 weeks ago

Selected Answer: D

Since need analyze the current EBS volume cost and to recommend optimizations, so have to use AWS compute optimizer
upvoted 2 times

Scheldon 1 year ago

Selected Answer: D

AnswerD

AWS Compute Optimizer helps avoid overprovisioning and underprovisioning four types of AWS resources—Amazon Elastic Compute Cloud (EC2) instance types, Amazon Elastic Block Store (EBS) volumes, Amazon Elastic Container Service (ECS) services on AWS Fargate, and AWS Lambda functions—based on your utilization data.

<https://aws.amazon.com/compute-optimizer/>

upvoted 4 times

Hkayne 1 year, 1 month ago

Selected Answer: D

Get recommendations to optimize your use of AWS resources
upvoted 3 times

Awsbeginner87 1 year, 2 months ago

Selected Answer: D

AWS Compute Optimizer helps avoid overprovisioning and underprovisioning four types of AWS resources—Amazon Elastic Compute Cloud (EC2) instance types, Amazon Elastic Block Store (EBS) volumes, Amazon Elastic Container Service (ECS) services on AWS Fargate, and AWS Lambda functions—based on your utilization data.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #868

Topic 1

A global company runs its workloads on AWS. The company's application uses Amazon S3 buckets across AWS Regions for sensitive data storage and analysis. The company stores millions of objects in multiple S3 buckets daily. The company wants to identify all S3 buckets that are not versioning-enabled.

Which solution will meet these requirements?

- B. Use Amazon S3 Storage Lens to identify all S3 buckets that are not versioning-enabled across Regions. **Most Voted**
- C. Enable IAM Access Analyzer for S3 to identify all S3 buckets that are not versioning-enabled across Regions.
- D. Create an S3 Multi-Region Access Point to identify all S3 buckets that are not versioning-enabled across Regions.

Correct Answer: B

Community vote distribution

B (100%)

Comments

joseantoniopolo **Highly Voted** 8 months ago

Selected Answer: B

<https://aws.amazon.com/blogs/aws/s3-storage-lens/>
upvoted 10 times

nadeerm **Most Recent** 3 months ago

Selected Answer: B

Answer is B
Option A: Set up an AWS CloudTrail event that has a rule to identify all S3 buckets that are not versioning-enabled across Regions.
upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: B

FYI, the missing Option A: Set up an AWS CloudTrail event that has a rule to identify all S3 buckets that are not versioning-enabled across Regions.

A - CloudTrail is meant for audit logs, not for analyzing bucket configurations directly.

B - Storage Lens can provide an overview of all S3 buckets and their versioning status across Regions. Those built-in features that can directly tackle your problems should be your top preferences.

C - IAM Access Analyzer is designed for access management, not storage configuration analysis.

D - S3 Multi-Region Access Points allow applications to access S3 buckets across Regions using a single global endpoint. So not relevant as well.

upvoted 3 times

sandordini 7 months, 1 week ago

Correct answer: A

A: You can use an AWS Config managed rule to identify Amazon S3 buckets that do not have versioning enabled.

upvoted 2 times

JA2018 5 months, 4 weeks ago

Potential Option A could "Use AWS Config rules to identify all S3 buckets that are not versioning-enabled across Regions "

upvoted 1 times

sandordini 7 months, 1 week ago

Question #889 contains all the answers including A, which is cloudtrail and obviously wrong.

On the other hand S3 Storage Lens:

You can use the Versioning-enabled bucket count metric to see which buckets use S3 Versioning. Then, you can take action in the S3 console to enable S3 Versioning for other buckets. So correct answer: B

upvoted 8 times

Awsbeginner87 8 months, 1 week ago

Where is option A

upvoted 2 times

xBUGx 8 months, 1 week ago

where is option A?

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #869

Topic 1

A company wants to enhance its ecommerce order-processing application that is deployed on AWS. The application must process each order exactly once without affecting the customer experience during unpredictable traffic surges.

Which solution will meet these requirements?

- A. Create an Amazon Simple Queue Service (Amazon SQS) FIFO queue. Put all the orders in the SQS queue. Configure an AWS Lambda function as the target to process the orders. **Most Voted**
- B. Create an Amazon Simple Notification Service (Amazon SNS) standard topic. Publish all the orders to the SNS standard topic. Configure the application as a notification target.
- C. Create a flow by using Amazon AppFlow. Send the orders to the flow. Configure an AWS Lambda function as the target to process the orders.
- D. Configure AWS X-Ray in the application to track the order requests. Configure the application to process the orders by pulling the orders from Amazon CloudWatch.

Correct Answer: A

Community vote distribution

A (100%)

Comments

mk168898 7 months, 2 weeks ago

Must process at least once => SQS
upvoted 2 times

MatAlves 8 months, 3 weeks ago

Selected Answer: A

"Standard queues support at-least-once message delivery, and FIFO queues support exactly-once message processing and high-throughput mode."

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/welcome.html#sqs-benefits>
upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: A

must process each order exactly once -> FIFO queue
upvoted 2 times

Scheldon 1 year ago

Selected Answer: A

AnswerA
SQS with FIFO queue will allow to read every customer order in order in which they came and only once.
upvoted 2 times

Hkayne 1 year, 1 month ago

Selected Answer: A

FIFO queue is the solution
upvoted 2 times

Tanidanindo 1 year, 2 months ago

Selected Answer: A

SQS and FIFO
upvoted 2 times

Kaula 1 year, 2 months ago

Selected Answer: A

FIFO > SQS
upvoted 2 times

Mikado211 1 year, 2 months ago

Selected Answer: A

The application must process each order exactly once == SQS + FIFO
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #870

Topic 1

A company has two AWS accounts: Production and Development. The company needs to push code changes in the Development account to the Production account. In the alpha phase, only two senior developers on the development team need access to the Production account. In the beta phase, more developers will need access to perform testing.

Which solution will meet these requirements?

- A. Create two policy documents by using the AWS Management Console in each account. Assign the policy to developers who need access.
- B. Create an IAM role in the Development account. Grant the IAM role access to the Production account. Allow developers to assume the role.
- C. Create an IAM role in the Production account. Define a trust policy that specifies the Development account. Allow developers to assume the role. **Most Voted**
- D. Create an IAM group in the Production account. Add the group as a principal in a trust policy that specifies the Production account. Add developers to the group.

Correct Answer: C*Community vote distribution***Comments****FlyingHawk** 4 months, 4 weeks ago**Selected Answer: C**

1. Creating an IAM role in the Production account with a trust relationship to the Development account.
 2. Grant the sts:AssumeRole permission to the specific developers (or groups) who need access to the Production account.
 3. During the beta phase, you can easily add more developers to the IAM group or policy in the Development account that allows them to assume the role in the Production account.
 4. In the Alpha phase, grant AssumeRole Permission to the groups of Two Senior Developers only
- upvoted 3 times

LeonSauveterre 5 months, 1 week ago**Selected Answer: C**

A - We need secure cross-account access between the Development and Production accounts. Option A did nothing about that.
B - This setup violates the AWS best practice for cross-account access, which recommends creating roles in the account being accessed (aka Production env).
C - As more developers need access, you can grant permissions in the Development account without modifying the role in Production. Also, access is granted through temporary credentials generated when the role is assumed, reducing the risk of long-term credential exposure.
D - Groups cannot establish trust between accounts. It didn't provide any mechanism for Development account users to access the Production account.

upvoted 2 times

LeonSauveterre 5 months, 1 week ago

To "assume" a role in AWS means that you (a developer) temporarily take on the permissions associated with that role. This is done using the AWS Security Token Service (STS), which generates temporary credentials.

Brief steps of option C:

1. Sets up a trust policy that allows entities from the Development account to assume the role. Attaches permissions for the required resources (in this case, access to specific services).
2. Attach an IAM policy to your user or role, allowing you to call "sts:AssumeRole".
3. Assume the role by doing this:

```
aws sts assume-role \  
--role-arn "arn:aws:iam::PRODUCTION_ACCOUNT_ID:role/ROLE_NAME" \  
--role-session-name "MySessionName"
```

4. Use the temporary credentials to access the resources in the Production account during the session duration.

upvoted 1 times

AMEJack 6 months, 2 weeks ago

Selected Answer: C

It should be C, groups can't be used in trust policy.

upvoted 2 times

Mayank0502 11 months, 1 week ago

Selected Answer: D

answer should be D

upvoted 1 times

f07ed8f 1 year ago

Selected Answer: C

https://docs.aws.amazon.com/IAM/latest/UserGuide/tutorial_cross-account-with-roles.html

upvoted 3 times

TwinSpark 1 year ago

Selected Answer: D

Weird question, but D is actually the only one that allow you to select which developer got access and when, so will go for D

upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Agree, as C will let any developers assume the role without control

upvoted 1 times

KennethNg923 11 months, 3 weeks ago

I check here: https://docs.aws.amazon.com/IAM/latest/UserGuide/tutorial_cross-account-with-roles.html, and yes it should be use IAM role, I correct my choice to C

upvoted 4 times

03beafc 1 year, 1 month ago

Selected Answer: A

you can't assign groups as principals, b and c don't specify only the senior devs, a is the only one that works here

upvoted 1 times

03beafc 1 year, 1 month ago

edit, none of these answers are right....
upvoted 1 times

Mikado211 1 year, 1 month ago

Selected Answer: D

If you want ALL the developers to assume the role in the production, then C using a trust policy to assume the role in production is perfect BUT

You could allow users in development account to assume the role in production, but in the end you will maintain potentially a big trust policy depending of the total number of users.

Here you want only some developers to connect to the production (others will follow without knowing if they all can connect and without knowing the number) so managing a separate group will give you a little more maintenance but will allow you to have different rights between the users.

I'd say D
upvoted 1 times

802c4ff 1 year, 1 month ago

Selected Answer: C

https://docs.aws.amazon.com/IAM/latest/UserGuide/tutorial_cross-account-with-roles.html
upvoted 3 times

xBUGx 1 year, 2 months ago

Selected Answer: D

i think D is better
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #871

Topic 1

A company wants to restrict access to the content of its web application. The company needs to protect the content by using authorization techniques that are available on AWS. The company also wants to implement a serverless architecture for authorization and authentication that has low login latency.


The solution must integrate with the web application and serve web content globally. The application currently has a small user base, but the company expects the application's user base to increase.

Which solution will meet these requirements?

- A. Configure Amazon Cognito for authentication. Implement Lambda@Edge for authorization. Configure Amazon CloudFront to serve the web application globally. **Most Voted**
- B. Configure AWS Directory Service for Microsoft Active Directory for authentication. Implement AWS Lambda for authorization. Use an Application Load Balancer to serve the web application globally.
- C. Configure Amazon Cognito for authentication. Implement AWS Lambda for authorization. Use Amazon S3 Transfer Acceleration to serve the web application globally.
- D. Configure AWS Directory Service for Microsoft Active Directory for authentication. Implement Lambda@Edge for authorization. Use AWS Elastic Beanstalk to serve the web application globally.

Correct Answer: A

Community vote distribution



Comments

Salilgen 5 months ago

Selected Answer: A

<https://aws.amazon.com/blogs/networking-and-content-delivery/authorizationedge-how-to-use-lambdaedge-and-json-web-tokens-to-enhance-web-application-security/>
upvoted 1 times

KennethNg923 11 months, 3 weeks ago

Authn: Cognito

Authz: Lambda@Edge

Globally: Lambda@Edge + cloudfront

upvoted 3 times

Hkayne 1 year, 1 month ago

Selected Answer: A

Serve content globally means the use of Cloudfront

upvoted 4 times

Danges 1 year, 1 month ago

Selected Answer: A

Implementación a nivel global ==> AWS Cloud Front

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #872

Topic 1

A development team uses multiple AWS accounts for its development, staging, and production environments. Team members have been launching large Amazon EC2 instances that are underutilized. A solutions architect must prevent large instances from being launched in all accounts.

How can the solutions architect meet this requirement with the LEAST operational overhead?

- A. Update the IAM policies to deny the launch of large EC2 instances. Apply the policies to all users.
- B. Define a resource in AWS Resource Access Manager that prevents the launch of large EC2 instances.
- C. Create an IAM role in each account that denies the launch of large EC2 instances. Grant the developers IAM group access to the role.
- D. Create an organization in AWS Organizations in the management account with the default policy. Create a service control policy (SCP) that denies the launch of large EC2 instances, and apply it to the AWS accounts. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Hkayne **Highly Voted** 1 year, 1 month ago

Selected Answer: D

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps.html
upvoted 6 times

LeonSauveterre **Most Recent** 5 months, 1 week ago

Selected Answer: D

A - You would need to update and maintain separate IAM policies in each account, which is too much trouble.
B - AWS Resource Access Manager (RAM) is primarily for resource sharing and does not directly restrict resource launches.
C - IAM roles are better suited for granting access rather than imposing restrictions. You would also need to create and maintain these roles in every account, even more trouble than option A.
D - SCPs are powerful here because they apply at the root account level, meaning that even if a developer has direct IAM permissions to launch large instances, the SCP will override and prevent it.

Also, there's no need to create or manage multiple IAM policies or roles across accounts. Once the SCP is defined and applied,

it enforces the restriction automatically.

upvoted 1 times

744fdad 10 months, 1 week ago

why is it not A? If the goal is only to prevent launch of EC2s

upvoted 2 times

78b9037 5 months, 3 weeks ago

Updating IAM policies must be done separately in each account, and policies for each user must be maintained, which will cause high operation overhead, so D.

upvoted 1 times

example_ 11 months ago

Selected Answer: D

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps_examples.html

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #873

Topic 1

A company has migrated a fleet of hundreds of on-premises virtual machines (VMs) to Amazon EC2 instances. The instances run a diverse fleet of Windows Server versions along with several Linux distributions. The company wants a solution that will automate inventory and updates of the operating systems. The company also needs a summary of common vulnerabilities of each instance for regular monthly reviews.

What should a solutions architect recommend to meet these requirements?

- A. Set up AWS Systems Manager Patch Manager to manage all the EC2 instances. Configure AWS Security Hub to produce monthly reports.
- B. Set up AWS Systems Manager Patch Manager to manage all the EC2 instances. Deploy Amazon Inspector, and configure monthly reports. **Most Voted**
- C. Set up AWS Shield Advanced, and configure monthly reports. Deploy AWS Config to automate patch installations on the EC2 instances.
- D. Set up Amazon GuardDuty in the account to monitor all EC2 instances. Deploy AWS Config to automate patch installations on the EC2 instances.

Correct Answer: B*Community vote distribution*

B (94%)

A (6%)

Comments**TwinSpark** **Highly Voted** 1 year ago**Selected Answer: B**

AWS Security Hub performs security best practice checks, aggregates alerts, and enables automated remediation.

Amazon Inspector scan workload for vulnerabilities

Guardduty threat detection for malicious activity in all account

AWS Shield DDOS

Regarding security B and D can be right (maybe D a little too much). For patching B is the only valid option.

upvoted 6 times

mk168898 **Most Recent** 7 months, 2 weeks ago

Inspector is for checking vulnerabilities, so B

upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: B

Inspector for Common vulnerabilities so it is B

upvoted 2 times

maxhg 1 year, 1 month ago

inspector for instances and software vulnerabilities

upvoted 4 times

Hkayne 1 year, 1 month ago

Selected Answer: B

AWS Systems Manager Patch Manager to automate the process of installing security-related updates for both the operating system and applications.

Amazon Inspector for Automated and continual vulnerability management at scale

upvoted 2 times

Alagong 1 year, 2 months ago

Selected Answer: A

Create an Auto Scaling group and an ELB in the DR Region, configuring the DynamoDB table as a global table, and setting up DNS failover to the new ELB. This approach allows for quick failover since the infrastructure is already in place and only DNS needs to be updated to redirect traffic.

upvoted 1 times

Tanidanindo 1 year, 2 months ago

Selected Answer: B

Inspector for vulnerability scanning

upvoted 3 times

Awsbeginner87 1 year, 2 months ago

Selected Answer: B

Option b

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #874


Topic 1

A company hosts its application in the AWS Cloud. The application runs on Amazon EC2 instances in an Auto Scaling group behind an Elastic Load Balancing (ELB) load balancer. The application connects to an Amazon DynamoDB table.

For disaster recovery (DR) purposes, the company wants to ensure that the application is available from another AWS Region with minimal downtime.

Which solution will meet these requirements with the LEAST downtime?

- A. Create an Auto Scaling group and an ELB in the DR Region. Configure the DynamoDB table as a global table. Configure DNS failover to point to the new DR Region's ELB. **Most Voted**
- B. Create an AWS CloudFormation template to create EC2 instances, ELBs, and DynamoDB tables to be launched when necessary. Configure DNS failover to point to the new DR Region's ELB.
- C. Create an AWS CloudFormation template to create EC2 instances and an ELB to be launched when necessary. Configure the DynamoDB table as a global table. Configure DNS failover to point to the new DR Region's ELB.
- D. Create an Auto Scaling group and an ELB in the DR Region. Configure the DynamoDB table as a global table. Create an Amazon CloudWatch alarm with an evaluation period of 10 minutes to invoke an AWS Lambda function that updates Amazon Route 53 to point to the DR Region's ELB.

Correct Answer: A*Community vote distribution***Comments****LeonSauveterre** 5 months, 1 week ago**Selected Answer: A**

A - By creating an Auto Scaling group and ELB in the DR Region upfront, you ensure that the infrastructure is ready to handle traffic immediately in case of a failure in the primary Region. This minimizes downtime during a disaster.

B & C - Waiting for the CloudFormation stack to create instances, ELBs, and DynamoDB tables delays recovery.

D - CloudWatch alarms and Lambda functions will introduce additional complexity and potential delays due to the evaluation period (≥ 10 minutes).

For A, simply create an Auto Scaling group and ELB in the DR Region. Use the same configurations as the primary Region. Just make sure that the AMI used for EC2 instances includes all necessary application code and dependencies.

make sure that the AMI used for EC2 instances includes all necessary application code and dependencies.

upvoted 2 times

mk168898 7 months, 2 weeks ago

B and C seems to take a while to launch, so it is not minimal down time

upvoted 2 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: A

Downtime I would choose Auto scaling + DNS failover rather than use cloud formation create infrastructure in DR region or Auto scaling + Lambda

upvoted 4 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: A

Downtime I would choose Auto scaling + DNS failover rather than use cloud formation create infrastructure in DR region or Auto scaling + Lambda

upvoted 3 times

Scheldon 1 year ago

Selected Answer: A

AnswerA

Hence there is no information that solution need to be cost effective and the main requirement is minimal downtime i would go with AUTOSCALING in DR region with 1 ELB and 1 server there but in case of need amount of servers can be increased automatically. Hence at least 1 server and ELB will be waiting and DynamoDB thanks to global table will be active in the same DR region as well, hence we need to inform users, using DNS about new destination, hence DNS failover to the ELB in DR region is the best solution here.

upvoted 4 times

Hkayne 1 year, 1 month ago

Selected Answer: A

With dynamo global tables, we just need to create an ELB and a ASG in the DR region resources. This resources will be used only if the main region fail over.

upvoted 2 times

Alagong 1 year, 2 months ago

Selected Answer: A

Create an Auto Scaling group and an ELB in the DR Region, configuring the DynamoDB table as a global table, and setting up DNS failover to the new ELB. This approach allows for quick failover since the infrastructure is already in place and only DNS needs to be updated to redirect traffic.

upvoted 4 times

Tanidanindo 1 year, 2 months ago

Selected Answer: A

Least downtime. C does not offer minimal downtime

upvoted 3 times

Awsbeginner87 1 year, 2 months ago

Selected Answer: C

Option C

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #875

Topic 1


A company runs an application on Amazon EC2 instances in a private subnet. The application needs to store and retrieve data in Amazon S3 buckets. According to regulatory requirements, the data must not travel across the public internet.

What should a solutions architect do to meet these requirements MOST cost-effectively?

- A. Deploy a NAT gateway to access the S3 buckets.
- B. Deploy AWS Storage Gateway to access the S3 buckets.
- C. Deploy an S3 interface endpoint to access the S3 buckets.
- D. Deploy an S3 gateway endpoint to access the S3 buckets. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

LeonSauveterre 5 months, 1 week ago

Selected Answer: D

About option C - S3 Interface Endpoint

1. Creates a private network interface (ENI) in your VPC that connects to Amazon S3.
2. Provides granular IAM-based controls at the operation level.
3. Typically much more expensive due to the cost associated with the network interface and data processing charges.

About option D - S3 Gateway Endpoint:

1. Creates a route entry in the VPC route table for accessing S3 without crossing the public internet.
2. Simpler and more cost-effective for high-volume S3 access, as there are no additional data processing charges or interface costs.
3. Doesn't support granular IAM controls per operation, but bucket policies can be used for access control so that's good.

upvoted 1 times

JA2018 5 months, 4 weeks ago

Selected Answer: C

Why I opt for Option C?

#Private Access:

It allows direct access to S3 buckets from a private subnet without data traversing the public internet, meeting the regulatory

It allows direct access to S3 buckets from a private subnet without data traversing the public internet, meeting the regulatory requirement.

#Cost-Effective:

Compared to other options, utilizing an S3 interface endpoint is generally the most cost-efficient way to achieve private access to S3.

#Summary:

When data needs to stay within a private network, use an S3 interface endpoint to access S3 buckets. Always consider the specific regulatory requirements and cost implications when choosing a solution.

upvoted 1 times

mk168898 7 months, 2 weeks ago

no internet and need access to s3 => s3 gateway endpoint

upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: D

Gateway endpoint free, so definitely interface end point expensive than it

upvoted 4 times

Scheldon 1 year ago

Selected Answer: D

AnswerD

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/privatelink-interface-endpoints.html>

taking into consideration that in both cases (s3 Instance Endpoint and S3Gateway endpoint), network traffic remains on the AWS network we need to think about other data which we have. For example application is in AWS cloud hence there is no need for access from on-premises. in that situation S3 Gateway endpoint seems to be better (and it is for free)

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/privatelink-interface-endpoints.html#types-of-vpc-endpoints-for-s3>

upvoted 4 times

Hkayne 1 year, 1 month ago

Selected Answer: D

D for sure.

upvoted 2 times

BatVanyo 1 year, 1 month ago

Selected Answer: D

Gateway endpoints are free.

upvoted 3 times

awsshare 1 year, 1 month ago

Selected Answer: D

Sorry, I think D is the correct option. Gateway endpoint is cheaper than Interface endpoint

upvoted 2 times

Tanidanindo 1 year, 2 months ago

Selected Answer: D

Gateway endpoint for S3

upvoted 4 times

awsshare 1 year, 2 months ago

Selected Answer: C

should be C

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #876

Topic 1

A company hosts an application on Amazon EC2 instances that run in a single Availability Zone. The application is accessible by using the transport layer of the Open Systems Interconnection (OSI) model. The company needs the application architecture to have high availability.

Which combination of steps will meet these requirements MOST cost-effectively? (Choose two.)

- A. Configure new EC2 instances in a different Availability Zone. Use Amazon Route 53 to route traffic to all instances.
- B. Configure a Network Load Balancer in front of the EC2 instances. **Most Voted**
- C. Configure a Network Load Balancer for TCP traffic to the instances. Configure an Application Load Balancer for HTTP and HTTPS traffic to the instances.
- D. Create an Auto Scaling group for the EC2 instances. Configure the Auto Scaling group to use multiple Availability Zones. Configure the Auto Scaling group to run application health checks on the instances. **Most Voted**
- E. Create an Amazon CloudWatch alarm. Configure the alarm to restart EC2 instances that transition to a stopped state.

Correct Answer: BD

Community vote distribution

BD (100%)

Comments

Scheldon **Highly Voted** 1 year ago

Selected Answer: BD

AnswerBD

There is no information about the ports numbers and only that means that we need to use NLB, ALB is only for HTTP hence there is no point of using that solution as connection can be done via any port for exampl 4000

Autoscaling with instances in multiple AZ is the best solution. it will allow run new EC2 if it fail and in case if whole AZ will go down we will have 2nd one.

upvoted 7 times

LuongTo **Most Recent** 7 months, 1 week ago

C is more comprehensive than B. I voted CD

upvoted 2 times

mk168898 7 months, 2 weeks ago

accessible by transport layer means not by application layer means no need ALB

upvoted 2 times

AbhiBK 9 months, 1 week ago

Answer is A & D - By deploying EC2 instances in multiple Availability Zones (Option A), you ensure that your application remains available even if one Availability Zone experiences an outage. This setup provides redundancy and fault tolerance.

upvoted 1 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: BD

as it mention "The application is accessible by using the transport layer" which is TCP, so no more information or reason to use ALB as well, so I will go for B+D

upvoted 3 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: BD

as it mention "The application is accessible by using the transport layer" which is TCP, so no more information or reason to use ALB as well, so I will go for B+D

upvoted 3 times

sandordini 1 year, 1 month ago

Selected Answer: BD

No word about the HTTP/application layer, only OSI 4 - TCP > B, an NLB should be enough
D: for Autoscaling.

upvoted 4 times

Tanidanindo 1 year, 2 months ago

Selected Answer: BD

transport layer means just NLB.

upvoted 3 times

Awsbeginner87 1 year, 2 months ago

Selected Answer: BD

B- since network layer operates at layer 4 i.e transport layer
D- for hHA

upvoted 2 times

Awsbeginner87 1 year, 2 months ago

Edited-D option for HA

upvoted 2 times

xBUGx 1 year, 2 months ago

Selected Answer: BD

question says the application is running on Transport Layer. i dont think there is need for ALB

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #877

Topic 1

A company uses Amazon S3 to host its static website. The company wants to add a contact form to the webpage. The contact form will have dynamic server-side components for users to input their name, email address, phone number, and user message.

The company expects fewer than 100 site visits each month. The contact form must notify the company by email when a customer fills out the form.

Which solution will meet these requirements MOST cost-effectively?

A. Host the dynamic contact form in Amazon Elastic Container Service (Amazon ECS). Set up Amazon Simple Email Service (Amazon SES) to connect to a third-party email provider.

B. Create an Amazon API Gateway endpoint that returns the contact form from an AWS Lambda function. Configure another Lambda function on the API Gateway to publish a message to an Amazon Simple Notification Service (Amazon SNS) topic.

Most Voted

C. Host the website by using AWS Amplify Hosting for static content and dynamic content. Use server-side scripting to build the contact form. Configure Amazon Simple Queue Service (Amazon SQS) to deliver the message to the company.

D. Migrate the website from Amazon S3 to Amazon EC2 instances that run Windows Server. Use Internet Information Services (IIS) for Windows Server to host the webpage. Use client-side scripting to build the contact form. Integrate the form with Amazon WorkMail.

Correct Answer: B

Community vote distribution

B (100%)

Comments

mk168898 7 months, 2 weeks ago

i see need to notify => SNS
upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: B

as it said "The company expects fewer than 100 site visits each month" and Lambda charge per each time calling it, so i will go

for B

upvoted 3 times

EdricHoang 11 months, 4 weeks ago

Selected Answer: B

Limitation is 100 submission a month, not frequent -> Lambda

upvoted 3 times

Scheldon 1 year ago

Selected Answer: B

AnswerB

API gateway + Lambda seems to be the best option expecially when SNS is in use which is specially created to push messaged to the subscribers (company appropraiate team in this situation)

upvoted 3 times

Hkayne 1 year, 1 month ago

Selected Answer: B

B is the right answer

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #878

Topic 1


A company creates dedicated AWS accounts in AWS Organizations for its business units. Recently, an important notification was sent to the root user email address of a business unit account instead of the assigned account owner. The company wants to ensure that all future notifications can be sent to different employees based on the notification categories of billing, operations, or security.

Which solution will meet these requirements MOST securely?

- A. Configure each AWS account to use a single email address that the company manages. Ensure that all account owners can access the email account to receive notifications. Configure alternate contacts for each AWS account with corresponding distribution lists for the billing team, the security team, and the operations team for each business unit.
- B. Configure each AWS account to use a different email distribution list for each business unit that the company manages. Configure each distribution list with administrator email addresses that can respond to alerts. Configure alternate contacts for each AWS account with corresponding distribution lists for the billing team, the security team, and the operations team for each business unit.
- C. Configure each AWS account root user email address to be the individual company managed email address of one person from each business unit. Configure alternate contacts for each AWS account with corresponding distribution lists for the billing team, the security team, and the operations team for each business unit.
- D. Configure each AWS account root user to use email aliases that go to a centralized mailbox. Configure alternate contacts for each account by using a single business managed email distribution list each for the billing team, the security team, and the operations team. **Most Voted**

Correct Answer: D

Community vote distribution

**Comments**

Linuslin **Highly Voted** 1 year ago

Selected Answer: D

D is correct answer.

Configuring each AWS account's root user to use email aliases that go to a centralized mailbox ensures that sensitive notifications are not directly sent to individual email addresses. Instead, they are directed to a centralized mailbox, reducing

the risk of unauthorized access to sensitive information. Additionally, configuring alternate contacts for each account using a single business-managed email distribution list for the billing team, the security team, and the operations team ensures that notifications are appropriately routed to the respective teams based on the categories of billing, operations, or security. This approach centralizes control and reduces the likelihood of misconfiguration or unauthorized access to sensitive notifications.

upvoted 12 times

Linuslin 1 year ago

Option A doesn't provide the same level of security because it relies on a single email address managed by the company, which could be compromised, and it requires all account owners to access this email account, potentially leading to access control issues.

Option B presents a similar issue as Option A, as it relies on different email distribution lists but still requires administrator email addresses that may be vulnerable to compromise.

Option C also has security concerns because it associates the root user email address with individual employees, which may not be ideal for security and access control purposes.

So, Option D is the most secure solution for ensuring that future notifications can be sent to different employees based on the notification categories of billing, operations, or security.

upvoted 10 times

d401c0d **Highly Voted** 1 year, 1 month ago

Selected Answer: A

Question mentions the email was sent to a business unit account instead of an account owner. Thus, A mentions all account owners to have access to email account.

upvoted 5 times

GOTJ **Most Recent** 3 months ago

Selected Answer: B

Option B is correct, according to AWS Prescriptive Guidance (see link below):

When setting up primary and alternate contacts for your AWS account, use an email distribution list instead of an individual's email address. Using an email distribution list makes sure that ownership and reachability are preserved as individuals in your organization come and go. Set alternate contacts for billing, operations, and security notifications, and use appropriate email distribution lists accordingly. AWS uses these email addresses to contact you, so it is important you retain access to them.

<https://docs.aws.amazon.com/prescriptive-guidance/latest/aws-startup-security-baseline/acct-01.html>

upvoted 2 times

bujuman 6 months, 3 weeks ago

Selected Answer: D

The most secure way should use "Update the alternate contacts for any AWS account in your organization"

<https://docs.aws.amazon.com/accounts/latest/reference/manage-acct-update-contact-alternate.html>

upvoted 3 times

MrIndia2022 7 months ago

question- option D uses a single distribution list for all business units vs Option B which uses separate, independent distribution lists for each business unit. This involves more segregation of emails and hence is more secure. I am confused between option B and d. Pl help

upvoted 1 times

MatAlves 8 months, 2 weeks ago

Selected Answer: D

Both B and D use Distribution Lists as part of the of the solution:

B - "configure each DL with administrator email addresses that can respond to alerts"

D - "configure alternate contacts for each account by using single business managed DL each" per team.

I confess the wording isn't amazing, but between B and D, the latter is the one that properly addresses the issue involving root user email address.

upvoted 3 times

jadeconstancy 11 months, 3 weeks ago

Selected Answer: D

correct answer is D
upvoted 4 times

sheilawu 1 year ago

Selected Answer: A

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_accounts_update_primary_email.html
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #879

Topic 1

A company runs an ecommerce application on AWS. Amazon EC2 instances process purchases and store the purchase details in an Amazon Aurora PostgreSQL DB cluster.

Customers are experiencing application timeouts during times of peak usage. A solutions architect needs to rearchitect the application so that the application can scale to meet peak usage demands.

Which combination of actions will meet these requirements MOST cost-effectively? (Choose two.)

- A. Configure an Auto Scaling group of new EC2 instances to retry the purchases until the processing is complete. Update the applications to connect to the DB cluster by using Amazon RDS Proxy. **Most Voted**
- B. Configure the application to use an Amazon ElastiCache cluster in front of the Aurora PostgreSQL DB cluster.
- C. Update the application to send the purchase requests to an Amazon Simple Queue Service (Amazon SQS) queue. Configure an Auto Scaling group of new EC2 instances that read from the SQS queue. **Most Voted**
- D. Configure an AWS Lambda function to retry the ticket purchases until the processing is complete.
- E. Configure an Amazon API Gateway REST API with a usage plan.

Correct Answer: AC

Community vote distribution



Comments

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: AC

A) uses RDS Proxy which is mainly for connection pooling and availability issues. Proxy is for too many connections(, not for performance: read replicas, caching)

B is caching which is designed for solving read-issues. (Here we have timeouts, and connection issues.)

C: SQS is good method for decoupling.

upvoted 8 times

Abdullah_Cloud **Highly Voted** 1 year, 1 month ago

Selected Answer: BC

I think it's BC
upvoted 6 times

7dcef09 Most Recent 1 month, 3 weeks ago

Selected Answer: BC

BC are better answers.
B improves performance which mitigates the application timeouts during times of peak usage.
C decouples the workload.
upvoted 1 times

tch 2 months, 2 weeks ago

Selected Answer: BC

application timeouts... not database connection timeout
upvoted 1 times

nadeerm 3 months ago

Selected Answer: BC

Amazon ElastiCache (Option B): ElastiCache can cache frequently accessed data, reducing the load on the Aurora PostgreSQL database. This improves read performance and reduces latency, which helps handle peak traffic more efficiently.

Amazon SQS with Auto Scaling (Option C): By sending purchase requests to an SQS queue, the application can decouple the front-end from the back-end processing. An Auto Scaling group of EC2 instances can process messages from the queue, ensuring that the system scales dynamically based on demand. This approach ensures that purchase requests are processed asynchronously, reducing the likelihood of timeouts during peak usage.
upvoted 1 times

ieffiong 3 months, 1 week ago

Selected Answer: BC

Combining A&C just leads to redundant auto scaling groups especially when you have SQS configured already. The question talks about combination of actions.
upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: AC

A - RDS Proxy improves scalability and database connection pooling.
B - Only helpful for read-heavy workloads. This question indicates "purchase transactions" (writes), so the main bottleneck is still there.
C - SQS allows for better handling of bursts in traffic by queueing requests (which is also true for so many questions before #879).
D - Cost-prohibitive at scale due to high invocation rates. And lambda is not suitable for long-running operations like database transactions.
E - Irrelevant.
upvoted 2 times

GOTJ 3 months, 3 weeks ago

Even though your RDS proxy approach is correct, by choosing "AC" you are configuring two ASGs: one to process the purchase (option C) and one for... Retry the purchase (?) (option A) This is not cost effective

In my opinion, designing architectures specifically to "retry" transactions (options AD), specially with timeout issues, is something that everyone should avoid, as it can easily degenerate into the "snowball" effect
upvoted 2 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: AB

why not choose B? C is wired design.
upvoted 1 times

Cpso 6 months, 1 week ago

Selected Answer: AB

As a developer choose both A&C is wired? while we use SQS solution to buffer request before handle by EC2 fleet. why we need EC2 in front of SQS from A? what of thier propose

need EC2 in front of SQS from A.: what do they propose.

upvoted 2 times

XXXXXINN 7 months, 3 weeks ago

BC.

Not idea why retry helps in this scenario besides it adds more complexity into the current design and also doesn't resolve the availability issue...

upvoted 4 times

b3b5fdd 7 months, 4 weeks ago

Selected Answer: BC

B and C!

upvoted 1 times

MatAlves 8 months, 2 weeks ago

Selected Answer: BC

A - simply pointless.

B- You're already using SQS (C), so why using ec2 to "retry the purchase"? They will stay in the queue until the purchase is processed. Otherwise, they will simply return to the queue.

C - This decouples the application from direct database calls, allowing the processing of purchase requests to scale independently and manage load more effectively.

upvoted 1 times

pujithacg8 9 months, 3 weeks ago

when we have SQS in option C why do you have to retry it again

I think the answer is B and C

upvoted 4 times

EdricHoang 11 months, 4 weeks ago

Selected Answer: AC

Combine SQS and auto-scaling EC2:

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-using-sqs-queue.html>

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #880

Topic 1

A company that uses AWS Organizations runs 150 applications across 30 different AWS accounts. The company used AWS Cost and Usage Report to create a new report in the management account. The report is delivered to an Amazon S3 bucket that is replicated to a bucket in the data collection account.

The company's senior leadership wants to view a custom dashboard that provides NAT gateway costs each day starting at the beginning of the current month.

Which solution will meet these requirements?

- A. Share an Amazon QuickSight dashboard that includes the requested table visual. Configure QuickSight to use AWS DataSync to query the new report.
- B. Share an Amazon QuickSight dashboard that includes the requested table visual. Configure QuickSight to use Amazon Athena to query the new report. **Most Voted**
- C. Share an Amazon CloudWatch dashboard that includes the requested table visual. Configure CloudWatch to use AWS DataSync to query the new report.
- D. Share an Amazon CloudWatch dashboard that includes the requested table visual. Configure CloudWatch to use Amazon Athena to query the new report.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**mk168898** 7 months, 2 weeks ago

dashboard => quicksight, S3 query => athena
upvoted 2 times

KennethNg923 11 months, 3 weeks ago**Selected Answer: B**

QuickSight for dashboard and Athena for query each month so it is B
upvoted 3 times

sandordini 1 year, 1 month ago

Selected Answer: B

Senior Leadership, custom dashboard, visualization: Quicksight Dashboard
S3 query: Athena
upvoted 3 times

d401c0d 1 year, 1 month ago

Selected Answer: B

B. Share an Amazon QuickSight dashboard that includes the requested table visual. Configure QuickSight to use Amazon Athena to query the new report.

QuickSight works well with Athena and it can interact S3
upvoted 2 times

Mikado211 1 year, 2 months ago

Selected Answer: B

You definitely use Athena to request S3.
Both cloudwatch and quicksight can interact with S3.
Since we are taking about "The company's senior leadership" I'd tend to use quicksight for a better format.
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #881

Topic 1

A company is hosting a high-traffic static website on Amazon S3 with an Amazon CloudFront distribution that has a default TTL of 0 seconds. The company wants to implement caching to improve performance for the website. However, the company also wants to ensure that stale content is not served for more than a few minutes after a deployment.

Which combination of caching methods should a solutions architect implement to meet these requirements? (Choose two.)

A. Set the CloudFront default TTL to 2 minutes. **Most Voted**

B. Set a default TTL of 2 minutes on the S3 bucket.


C. Add a Cache-Control private directive to the objects in Amazon S3.

D. Create an AWS Lambda@Edge function to add an Expires header to HTTP responses. Configure the function to run on viewer response.

E. Add a Cache-Control max-age directive of 24 hours to the objects in Amazon S3. On deployment, create a CloudFront invalidation to clear any changed files from edge caches. **Most Voted**

Correct Answer: AE

Community vote distribution



Comments

BBR01 **Highly Voted** 1 year, 1 month ago

Selected Answer: AE

AE.

By default, each file automatically expires after 24 hours, but you can change the default behavior in two ways:

1. To change the cache duration for all files that match the same path pattern, you can change the CloudFront settings for Minimum TTL, Maximum TTL, and Default TTL for a cache behavior.

2. To change the cache duration for an individual file, you can configure your origin to add a Cache-Control header with the max-age or s-maxage directive, or an Expires header to the file.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/Expiration.html#expiration-individual-objects>

upvoted 8 times

MatAlves **Highly Voted** 8 months, 2 weeks ago

Selected Answer: AE

Selected Answer: AC

You simply can't have A and E in the same approach:

"Default TTL applies only when your origin does not add HTTP headers such as Cache-Control max-age, Cache-Control s-maxage, or Expires to objects."

C - Cache-Control private directive specifies that the response is intended for a single user and should not be cached by shared caches - it can still be cached, but only on a client device.

This combination of steps would provide the best solution for the case.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/distribution-web-values-specify.html#DownloadDistValuesDefaultTTL>

upvoted 7 times

Sergantus 6 months, 3 weeks ago

Setting Cache-Control: private would prevent CloudFront from caching the content entirely, which is not the goal outlined, as it wants to use caching. After some time with updates, the caching performance will degrade for the entire solution as more and more objects get that directive.

upvoted 2 times

tch **Most Recent** 2 months, 2 weeks ago

Selected Answer: AE

Controlling how long Amazon S3 content is cached by Amazon CloudFront

<https://docs.aws.amazon.com/whitepapers/latest/build-static-websites-aws/controlling-how-long-amazon-s3-content-is-cached-by-amazon-cloudfront.html>

upvoted 1 times

tch 2 months, 2 weeks ago

Set maximum TTL value

Specify cache-control headers

upvoted 1 times

robotgeek 4 months, 3 weeks ago

Selected Answer: A

This question has gaps:

B) Can not be done in relation to S3 caching, this is a CloudFront specific configuration

C) Question does not make any comment in the sense of "private" objects... any assumption in that sense is your personal opinion, not the question stating it

D) Collisions with A, why would you do something like that programatically?

E) Can not be because why would you put 24h when it is not requested, also why would you invalidate a deployment that is already TTL=0, everything is up to date in the caches with that config

upvoted 1 times

Saliigen 5 months ago

Selected Answer: AC

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/Expiration.html#ExpirationDownloadDist>

upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: AE

A - Setting the default TTL to 2 minutes allows CloudFront to cache content for a short period.

B - TTL is not a concept applied directly to S3 buckets. TTL and caching behaviors are managed through Cache-Control headers and CloudFront settings.

C - The *private directive* lets caching occur locally only, not in shared caches like CloudFront.

D - Ensures content is served fresh by tightly controlling cache expiration, not even "a few minutes".

E - Using a long Cache-Control max-age directive improves performance, while targeted CloudFront invalidations ensure that updated content is served immediately after deployment. This is common practice.

upvoted 3 times

Cpso 6 months, 1 week ago

Selected Answer: AE

browser will always get cache-control value set by S3 not cloudfront.
cloudfront will use parameter ttl override the cache-control header for decision time to cache on cloudfront.

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/Expiration.html?utm_source=chatgpt.com#ExpirationDownloadDist
upvoted 2 times

Cpso 6 months, 1 week ago

Selected Answer: AE

A - make browser to cache 2 minutes. after 2 minute browser will ask cloudfront

E: cloudfront would not ask S3 for 24hour . invalidation will clearing this cacheing.

so A+E is most effective.

upvoted 2 times

JA2018 6 months, 1 week ago

Selected Answer: AD

#A - Set the CloudFront default TTL to 2 minutes: This directly controls how long content is cached at the CloudFront edge locations, allowing for a short caching window to meet the requirement of not serving stale content for more than a few minutes.

#D - Create a Lambda@Edge function to add an Expires header: By using a Lambda@Edge function triggered on the "viewer response", you can dynamically set an Expires header with a precise expiration time on each request, ensuring fine-grained control over caching behavior.

upvoted 1 times

JA2018 6 months, 1 week ago

#B - Set a default TTL of 2 minutes on the S3 bucket: This would not be as effective as setting the TTL at the CloudFront level as the caching would happen at the origin server (S3) instead of the edge locations, potentially impacting performance.

C - Add a Cache-Control private directive: This directive would prevent browser caching, which is not the desired outcome in this scenario.

E - Add a Cache-Control max-age directive of 24 hours with CloudFront invalidation: Setting a large max-age on objects in S3 and invalidating on deployment would not provide the necessary granular control for keeping content fresh within a few minutes.

upvoted 1 times

1ba9aa0 10 months, 1 week ago

Selected Answer: AC

A-C,

Because A-E is not possible following this link:

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/distribution-web-values-specify.html#DownloadDistValuesDefaultTTL>

"Default TTL applies only when your origin does not add HTTP headers such as Cache-Control max-age, Cache-Control s-maxage, or Expires to objects."

upvoted 3 times

EdricHoang 11 months ago

Selected Answer: AC

If the content still keep client's cache in 24h, its wrong (answer E)

upvoted 3 times

ug56c 12 months ago

Selected Answer: AE

If your minimum TTL is greater than 0, CloudFront uses the cache policy's minimum TTL, even if the Cache-Control: no-cache, no-store, and/or private directives are present in the origin headers.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/Expiration.html>

upvoted 5 times

Nm55560 1 year ago

nmk5509 1 year ago

Selected Answer: AE

"However, the company also wants to ensure that stale content is not served for more than a few minutes after a deployment."
After a deployment
upvoted 4 times

Sergantus 6 months, 3 weeks ago

Exactly, it was not outlined that the user shouldn't see the stale content, only that it's not served.
upvoted 1 times

Scheldon 1 year ago

Selected Answer: AC

Answer (AC)

Per table on URL

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/Expiration.html#expiration-individual-objects>
answer E is incorrect because if we will change cache-control max-age to 24h it will mean that customer browser will cache web for 24h and customer wants to be sure that it will be not longer than few min.
Expires header (answer D) from my understanding can be used only on full folder of web not as lambda function which will reply to customer requests.

We are setting Default TTL for CloudFront (answer A) not on S3 (answer B) and it will say CloudFront to cache web for 2min.
upvoted 4 times

Scheldon 1 year ago

Adding Cache-control private (answer C) will work per customer wish but only if we will add them to the objects which are changed very often or if we will set minimum TTL.

In the 1 situation User Browser will not store files which we designate to be often changed and mentioned files will be downloaded every time from origin.

In the 2 situation, Cloud front will cache web files for min TTL time but customer browser will not store them.

Taking all that in to account I would go with AC

upvoted 3 times

Linuslin 1 year ago

Selected Answer: AE

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/Expiration.html#expiration-individual-objects>

<https://stackoverflow.com/questions/43343759/confused-with-minimum-maximum-and-default-ttl-in-cloudfront>
upvoted 2 times

02ffe1c 1 year, 1 month ago

Selected Answer: DE

Since it don't want to cache more than a minute, A cannot be an answer
upvoted 1 times

mk168898 7 months, 2 weeks ago

no where in the question did it say 1 minute. You mean more than a few minutes?
upvoted 1 times

kelmryan1 1 year, 1 month ago

Answer is AE, C would only be on the user browser and would not cache to the cloud front and would be useless
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #882

Topic 1

A company runs its application by using Amazon EC2 instances and AWS Lambda functions. The EC2 instances run in private subnets of a VPC. The Lambda functions need direct network access to the EC2 instances for the application to work.


The application will run for 1 year. The number of Lambda functions that the application uses will increase during the 1-year period. The company must minimize costs on all application resources.

Which solution will meet these requirements?

- A. Purchase an EC2 Instance Savings Plan. Connect the Lambda functions to the private subnets that contain the EC2 instances.
- B. Purchase an EC2 Instance Savings Plan. Connect the Lambda functions to new public subnets in the same VPC where the EC2 instances run.
- C. Purchase a Compute Savings Plan. Connect the Lambda functions to the private subnets that contain the EC2 instances.
Most Voted
- D. Purchase a Compute Savings Plan. Keep the Lambda functions in the Lambda service VPC.

Correct Answer: C

Community vote distribution



Comments

Guru4Cloud **Highly Voted** 1 year, 1 month ago

Selected Answer: C

Compute Savings Plan: This plan offers significant discounts on Lambda functions compared to on-demand pricing. Since the application will run for a year, a sustained use discount like Compute Savings Plan is ideal.
Private Subnets: Lambda functions in private subnets can directly access EC2 instances within the VPC without needing internet access, reducing security risks and potential egress costs.

upvoted 7 times

0c21057 **Most Recent** 1 month, 3 weeks ago

Selected Answer: C

Compute Savings Plans provide the most flexibility and help to reduce your costs by up to 66%. These plans automatically

Compute Savings Plans provide the most flexibility and help to reduce your costs by up to 75% over those plans automatically apply to EC2 instance usage regardless of instance family, size, AZ, Region, OS or tenancy, and also apply to Fargate or Lambda usage. For example, with Compute Savings Plans, you can change from C4 to M5 instances, shift a workload from EU (Ireland) to EU (London), or move a workload from EC2 to Fargate or Lambda at any time and automatically continue to pay the Savings Plans price.

upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: C

Why Compute Savings Plan: Covers both EC2 and Lambda, providing flexibility as Lambda usage increases.

Why Private Subnets: Ensures network access between Lambda and EC2 instances.

!!

An EC2 Instance Savings Plan doesn't cover Lambda. Also, Lambda functions in public subnets cannot directly communicate with EC2 instances in private subnets.

upvoted 1 times

mk168898 7 months, 2 weeks ago

A and B is not because it is talkign about EC2, but the question is asking for Lambda

upvoted 2 times

jacinml 7 months, 3 weeks ago

Selected Answer: C

Compute savings include lamdba and EC2, Ec2 savings only EC2 instances.

<https://aws.amazon.com/savingsplans/compute-pricing/>

upvoted 2 times

Lin878 11 months, 1 week ago

I confuse this Question. Instance saving plan is cheaper than compute saving plan.

<https://aws.amazon.com/savingsplans/compute-pricing/>

upvoted 1 times

navneet_sh 11 months ago

But composite saving plan discount Automatically applies to Lambda.

upvoted 2 times

navneet_sh 11 months ago

Sorry I mean Compute Savings Plan.

upvoted 2 times

sheilawu 1 year ago

Selected Answer: A

In this question has point out "access EC2 instances" within VPC,=> Lambda VPC to an ENI (Elastic network interface) in your account VPC.=>No charge.

Therefore I stick with A, Not D.

upvoted 2 times

JohnYu 7 months, 3 weeks ago

An EC2 Instance Savings Plan is limited to savings on EC2 instances only. It does not provide cost savings for Lambda functions. Therefore, it is not the best choice given that the number of Lambda functions is expected to increase.

upvoted 3 times

sheilawu 1 year ago

I am sorry I mean C

upvoted 3 times

Scheldon 1 year ago

Selected Answer: D

AnswerD

<https://docs.aws.amazon.com/lambda/latest/dg/foundation-networking.html>
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #883

Topic 1

A company has deployed a multi-account strategy on AWS by using AWS Control Tower. The company has provided individual AWS accounts to each of its developers. The company wants to implement controls to limit AWS resource costs that the developers incur.

Which solution will meet these requirements with the LEAST operational overhead?

A. Instruct each developer to tag all their resources with a tag that has a key of CostCenter and a value of the developer's name. Use the required-tags AWS Config managed rule to check for the tag. Create an AWS Lambda function to terminate resources that do not have the tag. Configure AWS Cost Explorer to send a daily report to each developer to monitor their spending.


B. Use AWS Budgets to establish budgets for each developer account. Set up budget alerts for actual and forecast values to notify developers when they exceed or expect to exceed their assigned budget. Use AWS Budgets actions to apply a DenyAll policy to the developer's IAM role to prevent additional resources from being launched when the assigned budget is reached. **Most Voted**

C. Use AWS Cost Explorer to monitor and report on costs for each developer account. Configure Cost Explorer to send a daily report to each developer to monitor their spending. Use AWS Cost Anomaly Detection to detect anomalous spending and provide alerts.

D. Use AWS Service Catalog to allow developers to launch resources within a limited cost range. Create AWS Lambda functions in each AWS account to stop running resources at the end of each work day. Configure the Lambda functions to resume the resources at the start of each work day.

Correct Answer: B

Community vote distribution



Comments

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: B

My first instinct says B, but I'm concerned about the central management abilities of AWS Budgets. It seems that even though it is not planned to be used primarily to control other accounts it's still possible:

"You can use actions to define an explicit response that you want to take when a budget exceeds its action threshold. You can

trigger these alerts on actual or forecasted cost and usage budgets.

1. The management account sets the budget and threshold for the member account using budget filters.
2. When the budget threshold is breached, a budget action applies a restrictive SCP on the OU.

So hopefully B :D

upvoted 5 times

Scheldon Highly Voted 1 year ago

Selected Answer: B

AnswerB

<https://docs.aws.amazon.com/cost-management/latest/userguide/budgets-controls.html>

Taking into consideration that AWS Budgets is allowing to will inform you that you exceeded budgeted and execute actions like for example IAM actions to prevent running new resources in cloud, I think this option is a good and reasonable move. In case of need budgeted can be always increased and "chains" disabled.

upvoted 5 times

LeonSauveterre Most Recent 5 months, 1 week ago

Selected Answer: B

A - Tags help monitor costs, and AWS Config can enforce tagging rules, but manually tagging is TOO MUCH WORK.

B - Preventive (aka not reactive) cost control, good. Too provocative? Yeah, maybe. Developers might experience disruption if they hit their budget limits unexpectedly, and then the cost a little over the budget might not seem like a problem anymore... BUT this is actually a preventive approach with "the LEAST operational overhead".

C - "Daily reports" might be too late.

D - It can limit resource usage during non-working hours, but overall, you gotta deal with high operational overhead.

upvoted 2 times

LeonSauveterre 5 months, 1 week ago

The gist is, when you reach the limit, you must be stopped. It's the only way to prevent unexpected expenses. Reports, monitors, analyses, logs, charts, they DON'T work.

upvoted 1 times

MatAlves 8 months, 2 weeks ago

Selected Answer: B

As beautifully explained in this article:

<https://docs.aws.amazon.com/cost-management/latest/userguide/budgets-controls.html>

upvoted 2 times

NSA_Poker 11 months, 3 weeks ago

Selected Answer: B

(A) is eliminated. 'send a daily report to each developer' can be ignored.

(C) is eliminated. 'detect anomalous spending' won't stop the spending.

(D) is eliminated. 'stop running resources at the end of each work day' won't stop developers from mining bitcoin (\$\$\$) the next day.

(B) is correct. 'actions to apply a DenyAll policy..' is the only solution that will 'implement controls to limit AWS resource costs that the developers incur.'

upvoted 3 times

MatAlves 8 months, 2 weeks ago

I'd definitely be mining bitcoin the next day...

upvoted 3 times

f07ed8f 1 year ago

Selected Answer: C

Seem AWS Budgets does not have DenyAll function but only

Apply a custom Deny IAM policy that restricts the ability for a user, group, or role to provision additional Amazon EC2 resources

upvoted 2 times

BBR01 1 year, 1 month ago

Selected Answer: C

B and D are too aggressive.

A - "Instruct each developer", nope, too much operational work.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #884

Topic 1

A solutions architect is designing a three-tier web application. The architecture consists of an internet-facing Application Load Balancer (ALB) and a web tier that is hosted on Amazon EC2 instances in private subnets. The application tier with the business logic runs on EC2 instances in private subnets. The database tier consists of Microsoft SQL Server that runs on EC2 instances in private subnets. Security is a high priority for the company.

Which combination of security group configurations should the solutions architect use? (Choose three.)

A. Configure the security group for the web tier to allow inbound HTTPS traffic from the security group for the ALB.

Most Voted

B. Configure the security group for the web tier to allow outbound HTTPS traffic to 0.0.0.0/0.

C. Configure the security group for the database tier to allow inbound Microsoft SQL Server traffic from the security group for the application tier. Most Voted

D. Configure the security group for the database tier to allow outbound HTTPS traffic and Microsoft SQL Server traffic to the security group for the web tier.

E. Configure the security group for the application tier to allow inbound HTTPS traffic from the security group for the web tier. Most Voted

F. Configure the security group for the application tier to allow outbound HTTPS traffic and Microsoft SQL Server traffic to the security group for the web tier.

Correct Answer: ACE

Community vote distribution

ACE (100%)

Comments

EdricHoang Highly Voted 11 months, 4 weeks ago

Selected Answer: ACE

Security group is stateful, just need allow Inbound.
upvoted 6 times

KennethNg923 11 months, 3 weeks ago

Agree since security group is stateful so allow inbound is enough
upvoted 2 times

sk1974 Most Recent 3 months, 1 week ago

Selected Answer: ACE

Since SG's are stateful , allow inbound only. Ignore all outbound roles
upvoted 1 times

Scheldon 1 year ago

Selected Answer: ACE

AnswerACE:

Security Group is protecting instances, it's statefull. by default is allowing for outgoing traffic but not incomming.
hence we need to allow for inboud traffic. path looks like below
ALB >>HTTPS>> WEB tier >>HTTPS>> Application >>SQL traffic>> SQL DB
hence we need allow for
incoming https traffic on web tier
then
incomming http on app tier
and on the end for
incomming sql traffic on DB tier
upvoted 3 times

sandordini 1 year, 1 month ago

Selected Answer: ACE

ALB >>HTTPS>> WEB tier >>HTTPS>> Application >>SQL traffic>> SQL DB
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #885

Topic 1

A company has released a new version of its production application. The company's workload uses Amazon EC2, AWS Lambda, AWS Fargate, and Amazon SageMaker.


The company wants to cost optimize the workload now that usage is at a steady state. The company wants to cover the most services with the fewest savings plans.

Which combination of savings plans will meet these requirements? (Choose two.)

- A. Purchase an EC2 Instance Savings Plan for Amazon EC2 and SageMaker.
- B. Purchase a Compute Savings Plan for Amazon EC2, Lambda, and SageMaker.
- C. Purchase a SageMaker Savings Plan. **Most Voted**
- D. Purchase a Compute Savings Plan for Lambda, Fargate, and Amazon EC2. **Most Voted**
- E. Purchase an EC2 Instance Savings Plan for Amazon EC2 and Fargate.

Correct Answer: CD

Community vote distribution



Comments

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: CD

It's pretty obvious, although it's called: Machine Learning Savings Plans for Amazon SageMaker (C)
For the compute workloads we need a compute savings plan, that covers all the 3 compute options we use here (EC2, Lambda and Fargate) (D)
upvoted 8 times

mk168898 **Most Recent** 7 months, 2 weeks ago

just pick the 2 options that doesn't overlap
upvoted 4 times

mattyu 12 months ago

Selected Answer: CD

no doubt
upvoted 3 times

Scheldon 1 year ago

Answer CD

<https://aws.amazon.com/savingsplans/ml-pricing/>
<https://aws.amazon.com/savingsplans/compute-pricing/>
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #886

Topic 1

A company uses a Microsoft SQL Server database. The company's applications are connected to the database. The company wants to migrate to an Amazon Aurora PostgreSQL database with minimal changes to the application code.

Which combination of steps will meet these requirements? (Choose two.)

A. Use the AWS Schema Conversion Tool (AWS SCT) to rewrite the SQL queries in the applications.

B. Enable Babelfish on Aurora PostgreSQL to run the SQL queries from the applications. **Most Voted**

C. Migrate the database schema and data by using the AWS Schema Conversion Tool (AWS SCT) and AWS Database Migration Service (AWS DMS). **Most Voted**

D. Use Amazon RDS Proxy to connect the applications to Aurora PostgreSQL.

E. Use AWS Database Migration Service (AWS DMS) to rewrite the SQL queries in the applications.

Correct Answer: BC

Community vote distribution

BC (100%)

Comments

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: BC

B: Babelfish for Aurora PostgreSQL is a new capability for Amazon Aurora PostgreSQL-Compatible Edition that enables Aurora to understand commands from applications written for Microsoft SQL Server.

C: Is just obvious: Use Data Migration Tool for the migration, Schema Conversion tool for the Schema conversion.

upvoted 6 times

pranavsharma1604 1 year ago

<https://aws.amazon.com/rds/aurora/babelfish/>

upvoted 3 times

LeonSauveterre **Most Recent** 5 months, 1 week ago

Selected Answer: BC

A - AWS SCT does not directly rewrite SQL queries in the application code. This would require manual updates.

B - Babelfish allows Aurora PostgreSQL to understand SQL Server's native T-SQL language and protocol.

B - Babelfish allows Aurora PostgreSQL to understand SQL Server's native T-SQL language and protocol.

C - Convert the schema & migrate the data effectively, good.

D - Amazon RDS Proxy improves database scalability and availability (if the question mentions something like having trouble connecting during peak hours or whatnot then RDS Proxy is likely the answer) but does not address the compatibility issues between SQL Server and Aurora PostgreSQL.

E - This is not enough. Refer to option C.

upvoted 2 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: BC

DMS + SCT is correct, but "rewrite the SQL queries in the applications." is wrong so A + E are out.

Then only left B + C -> DMS + SCT + Babelfish (for SQL Server)

upvoted 2 times

Scheldon 1 year ago

Selected Answer: BC

AnswerBC

DMS will allow for DATABASE migration and use AWS Schema Conversion Tool (AWS SCT) to create some or all of the target tables, indexes, views, triggers, and so on.

<https://docs.aws.amazon.com/dms/latest/userguide/Welcome.html>

To minimize amount of code which need to be changed we need to use babelfish

<https://aws.amazon.com/rds/aurora/babelfish/>

upvoted 2 times

pranavsharma1604 1 year ago

Selected Answer: BC

<https://aws.amazon.com/rds/aurora/babelfish/>

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #887


Topic 1

A company plans to rehost an application to Amazon EC2 instances that use Amazon Elastic Block Store (Amazon EBS) as the attached storage.

A solutions architect must design a solution to ensure that all newly created Amazon EBS volumes are encrypted by default. The solution must also prevent the creation of unencrypted EBS volumes.

Which solution will meet these requirements?

- A. Configure the EC2 account attributes to always encrypt new EBS volumes. **Most Voted**
- B. Use AWS Config. Configure the encrypted-volumes identifier. Apply the default AWS Key Management Service (AWS KMS) key.
- C. Configure AWS Systems Manager to create encrypted copies of the EBS volumes. Reconfigure the EC2 instances to use the encrypted volumes.
- D. Create a customer managed key in AWS Key Management Service (AWS KMS). Configure AWS Migration Hub to use the key when the company migrates workloads.

Correct Answer: A*Community vote distribution***Comments****Scheldon** **Highly Voted** 1 year ago**Selected Answer: A**

AnswerA

The task is to force automatic encryption for every new EBS volume and prevent possibility of creation any unencrypted volume hence:

https://docs.aws.amazon.com/ebs/latest/userguide/work-with-ebs-encr.html#ebs-encryption_key_mgmt

To enable encryption by default for a Region

Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.

From the navigation bar, select the Region.

From the navigation pane, select EC2 Dashboard.

In the upper-right corner of the page, choose Account Attributes, Data protection and security.

Choose Manage.

Select Enable. You keep the AWS managed key with the alias alias/aws/ebs created on your behalf as the default encryption key, or choose a symmetric customer managed encryption key.

Choose Update EBS encryption.

upvoted 10 times

FlyingHawk Most Recent 4 months, 3 weeks ago

Selected Answer: A

<https://docs.aws.amazon.com/ebs/latest/userguide/encryption-by-default.html>

You can configure your AWS account to enforce the encryption of the new EBS volumes and snapshot copies that you create. For example, Amazon EBS encrypts the EBS volumes created when you launch an instance and the snapshots that you copy from an unencrypted snapshot. For examples of transitioning from unencrypted to encrypted EBS resources, see

<https://docs.aws.amazon.com/ebs/latest/userguide/ebs-encryption.html#encrypt-unencrypted>: To enable encryption by default for a Region

%aws ec2 enable-ebs-encryption-by-default --region region

In the upper-right corner of the page, choose Account Attributes, Data protection and security.

upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: B

A - This automatically encrypts all newly created volumes, but can't prevent unencrypted volumes.

B - AWS Config can enforce and automate compliance by using a rule to check whether all EBS volumes are encrypted.

C - Creating encrypted copies of volumes after they are created? Why bother to do this

D - What does Migration Hub have to do with this... I don't understand.

upvoted 2 times

Xandero 6 months ago

Selected Answer: A

<https://docs.aws.amazon.com/ebs/latest/userguide/encryption-by-default.html>

upvoted 1 times

EdricHoang 11 months ago

Selected Answer: B

"The solution must also prevent the creation of unencrypted EBS volumes."

For prevention future actions, I go for AWS config. You can setup Encryption in EC2, but Its manual process, what happen if you add one or more EC2?

upvoted 3 times

Scheldon 1 year ago

AnswerA

https://docs.aws.amazon.com/ebs/latest/userguide/work-with-ebs-encr.html#ebs-encryption_key_mgmt

To enable encryption by default for a Region

Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.

From the navigation bar, select the Region.

From the navigation pane, select EC2 Dashboard.

In the upper-right corner of the page, choose Account Attributes, Data protection and security.

Choose Manage.

Select Enable. You keep the AWS managed key with the alias alias/aws/ebs created on your behalf as the default encryption key, or choose a symmetric customer managed encryption key.

Choose Update EBS encryption

Choose update EBS encryption.

upvoted 1 times

Obdf3af 1 year ago

A. <https://repost.aws/knowledge-center/ebs-automatic-encryption>

upvoted 2 times

Isomas 1 year ago

Selected Answer: B

As it needs to prevent creation of Unencrypted EBS volume

upvoted 3 times

viejito 1 year ago

B es correcto , AWS Config para identificar automáticamente los volúmenes de EBS no cifrados y aplicar una acción correctiva.A,C,D : incorrectas , no cumplen con el cifrado automático

upvoted 3 times

JA2018 6 months, 1 week ago

From Google Translate

B is OK, AWS Config to automatically identify unencrypted EBS volumes and apply corrective action. A,C,D: Incorrect, do not comply with

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #888

Topic 1

An ecommerce company wants to collect user clickstream data from the company's website for real-time analysis. The website experiences fluctuating traffic patterns throughout the day. The company needs a scalable solution that can adapt to varying levels of traffic.

Which solution will meet these requirements?

- A. Use a data stream in Amazon Kinesis Data Streams in on-demand mode to capture the clickstream data. Use AWS Lambda to process the data in real time. **Most Voted**
- B. Use Amazon Kinesis Data Firehose to capture the clickstream data. Use AWS Glue to process the data in real time.
- C. Use Amazon Kinesis Video Streams to capture the clickstream data. Use AWS Glue to process the data in real time.
- D. Use Amazon Managed Service for Apache Flink (previously known as Amazon Kinesis Data Analytics) to capture the clickstream data. Use AWS Lambda to process the data in real time.

Correct Answer: A*Community vote distribution***Comments****bora4motion** 1 month ago**Selected Answer: A**

realtime - data streams, near real time - firehose. a is the correct answer.
upvoted 1 times

LeonSauveterre 5 months, 1 week ago**Selected Answer: A**

A - Kinesis Data Streams in on-demand mode automatically scales based on traffic, making it an ideal choice for fluctuating traffic patterns.
B - Kinesis Data Firehose is used for NEAR real-time data delivery. Also, AWS Glue is better suited for ETL and batch processing rather than real-time analysis.
C - Kinesis Video Streams is designed for video or audio streaming data, not clickstream data.
D - You can process data, but can't capture raw data first.

Also, AWS Lambda provides real-time data processing without the need for managing infrastructure.

upvoted 1 times

mk168898 7 months, 2 weeks ago

click stream => kinesis data streams

upvoted 2 times

MatAlves 8 months, 2 weeks ago

B - C are out = Glue doesn't support real-time data processing.

D - Why would you use Kinesis data ANALYTICS to ingest clickstream data instead of Amazon Kinesis DATA STREAM?

upvoted 2 times

Salilgen 5 months ago

<https://docs.aws.amazon.com/glue/latest/dg/streaming-chapter.html>

upvoted 1 times

tch 3 months ago

this is AWS Glue Streaming... not AWS Glue

upvoted 1 times

EdricHoang 11 months, 3 weeks ago

Selected Answer: A

This one is very tricky, need to read the context carefully:

"The company needs a scalable solution that can adapt to varying levels of traffic" -> Both Firehose and Stream are scalable. But, Firehose is automatic where Stream is not. However, the question does NOT say it should be automatic and Glue is not support real-time analysis.

Thats why go for A.

B is very close to.

upvoted 3 times

Scheldon 1 year ago

Selected Answer: A

AnswerA

Apache Flink (previously known as Amazon Kinesis Data Analytics) seems to not allowing sent data directly to Lambda...

Glue is allowing to integrate data from couple of sources in to one.

Hence I think A is correct answer

<https://docs.aws.amazon.com/lambda/latest/dg/with-kinesis.html>

<https://aws.amazon.com/kinesis/data-streams/features/?nc=sn&loc=2>

upvoted 2 times

f07ed8f 1 year ago

Selected Answer: A

Seem AWS Glue does not support process data in real time. I vote for A

upvoted 2 times

f07ed8f 1 year ago

Selected Answer: B

Both Kinesis Data Streams and Firehose are scalable but Firehose offers automated scaling. I vote fore B

upvoted 1 times

sandordini 1 year, 1 month ago

Selected Answer: A

I think Apache Flink (previously known as Amazon Kinesis Data Analytics) would also be fine, but as here it wants to combine it with Lambda, I would rather opt for Kinesis Data Streams + Lambda, so A, because of the figure on this page:

<https://aws.amazon.com/kinesis/>

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #889

Topic 1

A global company runs its workloads on AWS. The company's application uses Amazon S3 buckets across AWS Regions for sensitive data storage and analysis. The company stores millions of objects in multiple S3 buckets daily. The company wants to identify all S3 buckets that are not versioning-enabled.

Which solution will meet these requirements?

- A. Set up an AWS CloudTrail event that has a rule to identify all S3 buckets that are not versioning-enabled across Regions.
- B. Use Amazon S3 Storage Lens to identify all S3 buckets that are not versioning-enabled across Regions. **Most Voted**
- C. Enable IAM Access Analyzer for S3 to identify all S3 buckets that are not versioning-enabled across Regions.
- D. Create an S3 Multi-Region Access Point to identify all S3 buckets that are not versioning-enabled across Regions.

Correct Answer: B

Community vote distribution

B (100%)

Comments

mk168898 7 months, 2 weeks ago

sotrage lens used to identify non versioning
upvoted 1 times

Scheldon 1 year ago

Selected Answer: B

AnswerB

S3 Storage Lens "can also identify buckets that aren't following data-protection best practices, such as using S3 Replication or S3 Versioning."

https://docs.aws.amazon.com/AmazonS3/latest/userguide/storage_lens_basics_metrics_recommendations.html

upvoted 3 times

sandordini 1 year, 1 month ago

Selected Answer: B

You can use the Versioning-enabled bucket count metric to see which buckets use S3 Versioning. Then, you can take action in the S3 console to enable S3 Versioning for other buckets.

upvoted 2 times

✓

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #890

Topic 1

A company needs to optimize its Amazon S3 storage costs for an application that generates many files that cannot be recreated. Each file is approximately 5 MB and is stored in Amazon S3 Standard storage.


The company must store the files for 4 years before the files can be deleted. The files must be immediately accessible. The files are frequently accessed in the first 30 days of object creation, but they are rarely accessed after the first 30 days.

Which solution will meet these requirements MOST cost-effectively?

- A. Create an S3 Lifecycle policy to move the files to S3 Glacier Instant Retrieval 30 days after object creation. Delete the files 4 years after object creation. **Most Voted**
- B. Create an S3 Lifecycle policy to move the files to S3 One Zone-Infrequent Access (S3 One Zone-IA) 30 days after object creation. Delete the files 4 years after object creation.
- C. Create an S3 Lifecycle policy to move the files to S3 Standard-Infrequent Access (S3 Standard-IA) 30 days after object creation. Delete the files 4 years after object creation.
- D. Create an S3 Lifecycle policy to move the files to S3 Standard-Infrequent Access (S3 Standard-IA) 30 days after object creation. Move the files to S3 Glacier Flexible Retrieval 4 years after object creation.

Correct Answer: A

Community vote distribution



Comments

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: A

Although it's not stated what is meant by 'rarely accessed', this scenario would primarily be a candidate for the Glacier Instant Retrieval tier as the storage price would be more than 3 times lower compared to Standard IA. In the specific case of files being more frequently retrieved than quarterly, it can qualify for consideration of Standard IA.

Actually, we don't have the required info, so we have to guess what they are thinking.. which is pretty lame, to be honest..

upvoted 14 times

sheilawu 12 months ago

CC BY-SA 4.0 licensed content. All rights reserved. © 2021 ExamTopics.com

S3 Glacier Instant Retrieval only retain the file for maximum 90 days, that is why A is not correct answer.
So C is right answer.

upvoted 2 times

Dantecito 3 months, 2 weeks ago

The 90-day minimum storage duration means that if you delete, overwrite, or transition an object before 90 days, AWS will charge you for the remaining unused storage days.
It's a bit confusing but it's not like you can only have the files for 90 days, all glaciers are intended to retain data for archival purposes and 4 years as the question says is archival.

upvoted 2 times

Edwars 11 months, 1 week ago

No, it isn't. It is just the opposite.

Objects archived to S3 Glacier Flexible Retrieval have a minimum of 90 days of storage. If an object is deleted, overwritten, or transitioned before 90 days, a pro-rated charge equal to the storage charge for the remaining days will be incurred.

<https://aws.amazon.com/s3/faqs/?nc=sn&loc=7>

upvoted 3 times

MatAlves 8 months, 2 weeks ago

Hell no.

S3 Glacier Instant Retrieval is designed for long-lived, rarely accessed data that is retained for months or years. Objects that are archived to S3 Glacier Instant Retrieval have a minimum of 90 days of storage

upvoted 2 times

tch Most Recent 3 months ago

Selected Answer: A

B is out!

store the files for 4 years before the files can be deleted - for long term we should use S3 Glacier Instant Retrieval.

files are frequently accessed in the first 30 days of object creation, it actually store in Amazon S3 Standard storage

upvoted 1 times

tch 3 months ago

Each file is approximately 5 MB and is stored in Amazon S3 Standard storage, meaning when file created .. files already store in Amazon S3 Standard storage.. move to move the files to S3 Glacier Instant Retrieval 30 days after object creation is fine...

upvoted 1 times

a8a1e0e 3 months, 3 weeks ago

Selected Answer: C

Option C is correct instead of A because S3 Standard-IA provides lower long-term storage costs with immediate access, whereas S3 Glacier Instant Retrieval has higher retrieval costs, making it less cost-effective for files that may still need to be accessed occasionally.

upvoted 1 times

dfgdsfgfdgreg 4 months, 2 weeks ago

Selected Answer: A

Glacier is much cheaper for storage than standard. The catch is you have to pay for 90 days of storage no matter what, even if the object is deleted. Here we will store data for 4 years so glacier is the cheapest solution.

upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: C

A - S3 Glacier Instant Retrieval is designed for archival storage where objects are infrequently accessed but still require millisecond retrieval, so yes actually A will work. It's just that Glacier storage costs more than S3 Standard-IA (given that after 30 days, the files are only RARELY accessed), so this is not the most cost-effective option.

B - For files that "cannot be recreated," durability is crucial, so single Availability Zone is not OK.

C - S3 Standard-IA is optimized for objects accessed less frequently but still requires rapid retrieval and high durability.

D - After 4 years, let's just delete those files.

upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Just a thought, if not constrained to these option, I would simply use Intelligent-Tiering to automatically move the files between access tiers (frequent access, infrequent access, etc) based on actual usage. Eventually, I'd delete these files after 4 years as well. I think this approach is better.

upvoted 1 times

MatAlves 8 months, 2 weeks ago

Selected Answer: A

Instant Retrieval = immediately accessible.

Choose the S3 Glacier Instant Retrieval storage class when you need milliseconds access to low cost archive data.

<https://aws.amazon.com/s3/faqs/?nc=sn&loc=7>

upvoted 3 times

scaredSquirrel 9 months ago

Selected Answer: A

Glacier Instant Retrieval is much cheaper, and it is intended for archival storage with very low access patterns

upvoted 3 times

Ucy 10 months, 2 weeks ago

Selected Answer: C

While S3 Glacier Instant Retrieval offers immediate access, it has a minimum storage duration policy, Objects stored in S3 Glacier Instant Retrieval have a minimum storage duration of 90 days.

upvoted 2 times

FrozenCarrot 11 months ago

Selected Answer: A

S3 Glacier Instant Retrieval delivers the fastest access to archive storage, with the same throughput and milliseconds access as the S3 Standard and S3 Standard-IA storage classes.

--- <https://aws.amazon.com/s3/storage-classes/glacier/instant-retrieval/>

upvoted 4 times

EdricHoang 11 months ago

Selected Answer: B

cost effective -> OneZone IA

"The files must be immediately accessible" -> cannot be Glacier

upvoted 1 times

MatAlves 8 months, 2 weeks ago

Instant Retrieval = immediately accessible.

Choose the S3 Glacier Instant Retrieval storage class when you need milliseconds access to low cost archive data.

<https://aws.amazon.com/s3/faqs/?nc=sn&loc=7>

upvoted 2 times

Johnoppong101 9 months, 3 weeks ago

files can not be recreated -> OneZone IA not accepted

upvoted 2 times

Mayank0502 11 months, 1 week ago

Selected Answer: C

this option has most durability

upvoted 1 times

Lin878 11 months, 1 week ago

Selected Answer: C

"are rarely accessed after the first 30 days" - not often
I will go with "C".

upvoted 1 times

345645a 11 months, 2 weeks ago

<https://aws.amazon.com/es/s3/storage-classes/glacier/instant-retrieval/>

upvoted 2 times

sheilawu 1 year ago

Selected Answer: C

Immediately accessible =>C

upvoted 1 times

Scheldon 1 year ago

Selected Answer: C

AnswerC

We cannot choose B because if that one zone will fail, company will not be able to recreate them.

We cannot choose D because we do not have to store files after 4y hence we can delete them (cost savings)

We cannot choose A - Glacier is less expensive (0,004 per GB) than S3-Standard - IA but is not allowing for instant access which is one of requirements (there is no information that data access shouldn't be accessible immediately) we have only information that after 30d access to data is less frequently. Hence I think we need to choose S3 Standard - IA (answer C)

upvoted 4 times

bujuman 1 year ago

Selected Answer: C

Requirements:

- frequently accessed for 30 days
- lower cost

Features for S3 Standard-IA:

- Infrequently accessed objects
- Milliseconds to access

According to me, best option for this use case is C

NB: Glacier better suits for lower cost, infrequent access.

upvoted 2 times

th1002 1 year ago

Selected Answer: C

why do we need one zone, glacier instant for 30 days ? or why do we need to move to glacier after 4 years ?

I think C is correct

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #891

Topic 1

A company runs its critical storage application in the AWS Cloud. The application uses Amazon S3 in two AWS Regions. The company wants the application to send remote user data to the nearest S3 bucket with no public network congestion. The company also wants the application to fail over with the least amount of management of Amazon S3.

Which solution will meet these requirements?

- A. Implement an active-active design between the two Regions. Configure the application to use the regional S3 endpoints closest to the user.
- B. Use an active-passive configuration with S3 Multi-Region Access Points. Create a global endpoint for each of the Regions.
- C. Send user data to the regional S3 endpoints closest to the user. Configure an S3 cross-account replication rule to keep the S3 buckets synchronized.
- D. Set up Amazon S3 to use Multi-Region Access Points in an active-active configuration with a single global endpoint. Configure S3 Cross-Region Replication. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: D

Using a Multi-region Accesspoint in an Active-Active setup will send data to the closest Region, without accessing the internet: "send remote user data to the nearest S3 bucket with no public network congestion"

Not very easy to read and understand but it's all there:

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/MultiRegionAccessPoints.html>

upvoted 16 times

LeonSauveterre **Most Recent** 5 months, 1 week ago

Selected Answer: D

A - The "implementation" is so much more complicated than you think.

B - Active-passive? Then you gotta promote the passive region to active manually in case of failure.

C - This option does nothing to achieve low-latency or failover requirements.

D - It simplifies configuration and management with a single global endpoint, and provides seamless failover and high availability through Multi-Region Access Points and Cross-Region Replication.

upvoted 1 times

LeonSauveterre 5 months, 1 week ago

And yes, the single global endpoint in option D will automatically route traffic to the nearest bucket, which ensures low-latency, which is a plus.

upvoted 1 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: D

D. Set up Amazon S3 to use Multi-Region Access Points in an active-active configuration with a single global endpoint. Configure S3 Cross-Region Replication.

upvoted 2 times

ike001 11 months, 3 weeks ago

D is correct

upvoted 3 times

Scheldon 1 year ago

Selected Answer: D

Answer D

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/MultiRegionAccessPoints.html>

When you create a Multi-Region Access Point, you specify a set of AWS Regions where you want to store data to be served through that Multi-Region Access Point. You can use S3 Cross-Region Replication (CRR) to synchronize data among buckets in those Regions. You can then request or write data through the Multi-Region Access Point global endpoint. Amazon S3 automatically serves requests to the replicated dataset from the closest available Region. Multi-Region Access Points are also compatible with applications that are running in Amazon virtual private clouds (VPCs), including those that are using AWS PrivateLink for Amazon S3.

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/replication.html>

upvoted 3 times

1223d0e 1 year, 1 month ago

To me it looks like C, the requirement is to send the request to the closest region

upvoted 1 times

trongod05 8 months ago

It's D. In an active-active configuration, requests made to an S3 Multi-Region Access Point's global endpoint automatically route over the AWS global network to the nearest S3 bucket. This allows applications to automatically avoid congested network segments on the public internet, improving application performance and reliability.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #892

Topic 1

A company is migrating a data center from its on-premises location to AWS. The company has several legacy applications that are hosted on individual virtual servers. Changes to the application designs cannot be made.

Each individual virtual server currently runs as its own EC2 instance. A solutions architect needs to ensure that the applications are reliable and fault tolerant after migration to AWS. The applications will run on Amazon EC2 instances.

Which solution will meet these requirements?

A. Create an Auto Scaling group that has a minimum of one and a maximum of one. Create an Amazon Machine Image (AMI) of each application instance. Use the AMI to create EC2 instances in the Auto Scaling group. Configure an Application Load Balancer in front of the Auto Scaling group. **Most Voted**


B. Use AWS Backup to create an hourly backup of the EC2 instance that hosts each application. Store the backup in Amazon S3 in a separate Availability Zone. Configure a disaster recovery process to restore the EC2 instance for each application from its most recent backup.

C. Create an Amazon Machine Image (AMI) of each application instance. Launch two new EC2 instances from the AMI. Place each EC2 instance in a separate Availability Zone. Configure a Network Load Balancer that has the EC2 instances as targets.

D. Use AWS Mitigation Hub Refactor Spaces to migrate each application off the EC2 instance. Break down functionality from each application into individual components. Host each application on Amazon Elastic Container Service (Amazon ECS) with an AWS Fargate launch type.

Correct Answer: A

Community vote distribution

**Comments**

Vasiliy **Highly Voted** 1 year ago

Selected Answer: A

Autoscaling with max=1 is what is needed to keep only one instance at a time - it will still fail, but it will spawn exactly one instance in case of failure (we are not allowed to change the design of the app)
Having single instances in different AZ will not help - if one of the AZs is down, the app will still be affected

upvoted 13 times

sandordini Highly Voted 1 year, 1 month ago

Selected Answer: C

NOT A: Autoscaling with Maximum of 1 EC2 :D
NOT B: Hourly backup... RPO 1hr
C: AMI, Multi-AZ -> Fault tolerant
NOT D: ECS with Fargate, but it needs to run on EC2..

upvoted 11 times

Scheldon 11 months, 1 week ago

You cannot change application and based on the story in the past there was only one server/copy of application running at the same time. So we cannot run more than one copy of any application at once.

It is possible to set Min and Max to 1 which will automatically bring up server when it will crash. Taking into consideration that we cannot change application design and load-balancing between regions would probably need that (no information if applications are stateful or stateless) I would go for solution in answer A

upvoted 5 times

bora4motion Most Recent 1 month ago

Selected Answer: C

needs to be C for fault tolerant.

upvoted 1 times

xerxusu 2 months, 2 weeks ago

Selected Answer: A

这个题，真反人类

upvoted 1 times

Dantecito 3 months, 2 weeks ago

Selected Answer: C

C is the answer.

A. Changes in application designs are not the same as changes in infrastructure, also is not fault tolerant and an ALB just for one instance? what's the point of that

upvoted 1 times

dfgdsfgdgreg 4 months, 2 weeks ago

Selected Answer: C

A has a single point of failure = not fault tolerant

upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: C

A - Creating an AMI of each application instance ensures that the migrated applications are deployed consistently with their existing configurations, yes, but an Auto Scaling group with a minimum and maximum of one provides no redundancy. If the single instance fails, the application becomes unavailable.

B - Backup-based disaster recovery does not provide real-time fault tolerance. It introduces downtime during recovery and does not meet the reliability requirement.

C - Launching two EC2 instances in separate Availability Zones (AZs) provides high availability. If one AZ experiences an outage, the application remains accessible on the instance in the other AZ.

D - "Changes to the application designs cannot be made."

upvoted 2 times

LeonSauveterre 5 months, 1 week ago

I copied this about option A. It's got some valid points, but just for your references.

An Application Load Balancer is also not the best fit for all legacy applications, as it operates at the application layer (of 7-layer model) and requires HTTP/HTTPS traffic. So for protocols like FTP, SMTP or other TCP-based customized protocols, an ALB would be unsuitable.

upvoted 1 times

IA2018 5 months, 1 week ago

JAZU18 6 months, 1 week ago

Selected Answer: C

I will choose Option C for the following reasons:

#1 AMI creation: Creating an AMI of each application instance allows for easy replication.

#2 Multiple Availability Zones: Launching instances in different Availability Zones provides redundancy in case of regional outages.

#3 Network Load Balancer: A Network Load Balancer distributes traffic across the instances, providing high availability by automatically routing requests to a healthy instance if one fails.

upvoted 2 times

daveya 6 months, 2 weeks ago

What if application is statefull and there's no backup?

upvoted 2 times

zits88 7 months, 2 weeks ago

Selected Answer: C

I don't see how people can choose A here when it is talking about reliability and the answer only has a maximum of 1. There's nothing about cost effectiveness so there is no reason this one is better. Definitely C.

upvoted 1 times

Cpso 6 months, 1 week ago

we cant change applicaiton logic to be run 2 instance concurrently.

upvoted 1 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: A

A. Create an Auto Scaling group that has a minimum of one and a maximum of one. Create an Amazon Machine Image (AMI) of each application instance. Use the AMI to create EC2 instances in the Auto Scaling group. Configure an Application Load Balancer in front of the Auto Scaling group.

upvoted 3 times

EdricHoang 11 months, 1 week ago

Selected Answer: C

Fault tolerance is not High Availability

Answer A is HA design, not Fault tolerance

upvoted 2 times

Cpso 6 months, 1 week ago

exam said we can change application logic. 99.9% of single app usually work strange if you run it 2 process concurrently and random each request to them.

upvoted 1 times

Lin878 11 months, 1 week ago

Selected Answer: A

A makes sense.

C is possible but manual intervention.

upvoted 2 times

emakid 11 months, 2 weeks ago

answer is A.

Auto Scaling Group:

Explanation:

Minimum and Maximum of one instance: Ensures that the instance is always running. If the instance fails, Auto Scaling will automatically replace it with a new one, maintaining high availability.

Amazon Machine Image (AMI): Captures the current state of the application instance, ensuring that new instances launched by Auto Scaling will have the same configuration.

Application Load Balancer (ALB):

Load Balancer: Distributes traffic to the instances in the Auto Scaling group, ensuring fault tolerance. Even though there is only one instance, the ALB can help manage incoming traffic and be ready for future scaling if needed.

For C:

While this provides high availability, it does not address fault tolerance as effectively as the Auto Scaling group approach. Without Auto Scaling, if an instance fails, manual intervention is required to launch new instances.

upvoted 4 times

24b2e9e 11 months, 2 weeks ago

A makes sense

upvoted 2 times

Nm55569 1 year ago

Selected Answer: A

It's either A or B but A is a better option. The application design cannot be changed so we don't know if it can run across 2 servers.

upvoted 3 times

Scheldon 1 year ago

Selected Answer: A

Answer A

It is possible to set Min and Max to 1 which will automatically bring up server when it will crash. Taking into consideration that we cannot change application design and load-balancing between regions would probably need that (no information if applications are statefull or stateless) i would go for solution in answer A

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/asg-capacity-limits.html>

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #893

Topic 1


A company wants to isolate its workloads by creating an AWS account for each workload. The company needs a solution that centrally manages networking components for the workloads. The solution also must create accounts with automatic security controls (guardrails).

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use AWS Control Tower to deploy accounts. Create a networking account that has a VPC with private subnets and public subnets. Use AWS Resource Access Manager (AWS RAM) to share the subnets with the workload accounts. **Most Voted**
- B. Use AWS Organizations to deploy accounts. Create a networking account that has a VPC with private subnets and public subnets. Use AWS Resource Access Manager (AWS RAM) to share the subnets with the workload accounts.
- C. Use AWS Control Tower to deploy accounts. Deploy a VPC in each workload account. Configure each VPC to route through an inspection VPC by using a transit gateway attachment.
- D. Use AWS Organizations to deploy accounts. Deploy a VPC in each workload account. Configure each VPC to route through an inspection VPC by using a transit gateway attachment.

Correct Answer: A

Community vote distribution



Comments

bujuman **Highly Voted** 1 year ago

Selected Answer: A

Statement:

- The solution also must create accounts with automatic security controls (guardrails).

<https://docs.aws.amazon.com/controltower/latest/userguide/what-is-control-tower.html>

AWS Control Tower provides a pre-packaged set of guardrails (policies) and blueprints (best-practice configurations) to ensure that the environment complies with security and compliance standards. It's designed to simplify the process of creating and managing a multi-account AWS environment while maintaining security and compliance.

upvoted 9 times

sandordini **Highly Voted** 1 year, 1 month ago

Selected Answer: B

It's a hard one. I'd go for B
Several accounts in an org, with central mgmt > AWS Organization
Sharing resources among accounts > AWS RAM
AWS Organizations and RAM typically work well together...
Happy to be challenged, of course.
upvoted 6 times

sandordini 1 year, 1 month ago

Although automatic security control could be a hint for AWS Control Tower
(set up and operate your multi-account AWS environment with prescriptive controls)
upvoted 1 times

FlyingHawk Most Recent 4 months, 2 weeks ago

Selected Answer: A

AWS Control Tower automates the setup of accounts and guardrails, reducing the need for manual configuration.

Centralizing networking in a single account and sharing resources via AWS RAM minimizes the operational effort required to manage networking across multiple accounts.

upvoted 1 times

LeonSauveterre 5 months, 1 week ago

Selected Answer: A

A - It works but on first sight I eliminated this option because if you have a large number of workloads, the manual work would be too much. However, other options have bigger issues so this one is pretty great actually.

B - AWS Organizations alone does not provide the automatic guardrails, so it requires manual implementation of security controls and policies.

C - Also a valid choice, especially if more isolation and advanced traffic inspection are needed, but it introduces more complexity than option A and might require more operational oversight, which would not be ideal for "least operational overhead."

D - Like B.

upvoted 1 times

mk168898 7 months, 2 weeks ago

guard rails => AWS control tower

upvoted 2 times

XXXXXINN 8 months ago

A
Guardrails >> AWS Control Tower

upvoted 3 times

dhewa 9 months, 3 weeks ago

Selected Answer: A

AWS Control Tower provides built-in guardrails and automates the creation of accounts with security controls.

upvoted 2 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: A

A. Use AWS Control Tower to deploy accounts. Create a networking account that has a VPC with private subnets and public subnets. Use AWS Resource Access Manager (AWS RAM) to share the subnets with the workload accounts.

upvoted 2 times

emakid 11 months, 2 weeks ago

Selected Answer: A

It leverages AWS Control Tower for automated account deployment and management, along with AWS RAM for centralized networking management, thus minimizing operational overhead while meeting the company's requirements for workload isolation and automatic security controls.

upvoted 3 times

stalk98 1 year ago

Selected Answer: A

answer is A
upvoted 2 times

Tomrr 1 year ago

Selected Answer: A

Anser is A, Control Tower has guardrails

AWS Audit Manager provides an AWS Control Tower Guardrails framework to assist you with your audit preparation.
upvoted 2 times

Scheldon 1 year ago

Selected Answer: A

Taking into consideration that AWS Control Tower is Orchestrator for AWS Organization which applies guardrails, I think A is a good choose.

<https://docs.aws.amazon.com/controltower/latest/userguide/what-is-control-tower.html>
upvoted 3 times

1223d0e 1 year, 1 month ago

Please explain why the answer is option A
upvoted 1 times

jackey_feng 1 year ago

I prefer B which is free. A may cause fee for sevice used while I am not sure about it.
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #894

Topic 1


A company hosts a website on Amazon EC2 instances behind an Application Load Balancer (ALB). The website serves static content. Website traffic is increasing. The company wants to minimize the website hosting costs.

Which solution will meet these requirements?

- A. Move the website to an Amazon S3 bucket. Configure an Amazon CloudFront distribution for the S3 bucket. **Most Voted**
- B. Move the website to an Amazon S3 bucket. Configure an Amazon ElastiCache cluster for the S3 bucket.
- C. Move the website to AWS Amplify. Configure an ALB to resolve to the Amplify website.
- D. Move the website to AWS Amplify. Configure EC2 instances to cache the website.

Correct Answer: A

Community vote distribution



Comments

Scheldon **Highly Voted** 1 year ago

Selected Answer: A

Answer A

Based on <https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/getting-started-secure-static-website-cloudformation-template.html>

Amazon CloudFront:

Uses the durable storage of Amazon Simple Storage Service (Amazon S3) – This solution creates an Amazon S3 bucket to host your static website's content. To update your website, just upload your new files to the S3 bucket.

upvoted 9 times

Dantecito **Most Recent** 3 months, 2 weeks ago

Selected Answer: A

I don't think anyone has mentioned it, but cloudfront helps reduce costs. If you want to use ElastiCache, it is possible but incredibly difficult, the question asks to reduce costs and cloudfront is cheaper than ElastiCache.

upvoted 1 times

joanna91 4 months, 3 weeks ago

Selected Answer: A

A vs B -> A. Because Elsticache is for DB, not for S3.

upvoted 1 times

Cpso 6 months, 1 week ago

Selected Answer: A

A. but questions should specify it is static or dynamic website

upvoted 1 times

j21kjf0o9ijhj 6 months, 4 weeks ago

Selected Answer: A

Why??? Guys, you'll fail the others. Don't troll or believe everything chatgpt says... ☐

The question indicates "The website serves static content", why would you use elasticache for it? It may work, but it'll take a helluva effort.

S3 and CloudFront is straightforward and simple. Not only that, it also works great for static content. It's AWS's cdn. It can also serve dynamic content, but that's out of scope for this question.

<https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/getting-started-secure-static-website-cloudformation-template.html>

upvoted 2 times

mk168898 7 months, 2 weeks ago

Selected Answer: B

S3 bucket is correct, so it's between A or B.

B because cache to handle when traffic increase.

upvoted 1 times

mk168898 7 months, 2 weeks ago

i changed my answer to A because elasticache is only for dynamic content but question only requires static content

upvoted 2 times

Abdullah2004 9 months, 2 weeks ago

Selected Answer: B

website serves static content So we need to ElastiCache when traffic increases , no need for cloudfront

upvoted 1 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: A

A. Move the website to an Amazon S3 bucket. Configure an Amazon CloudFront distribution for the S3 bucket.

upvoted 3 times

FrozenCarrot 11 months, 1 week ago

Selected Answer: B

S3 for static contents, and ElastiCache for traffic increasing. No need for cloudfront cuz there is no need for global deliver for the website

upvoted 3 times

FrozenCarrot 11 months ago

" customers often complement S3 with an in-memory cache, such as Amazon ElastiCache for Redis, to reduce the S3 retrieval cost and to improve performance." --<https://aws.amazon.com/blogs/storage/turbocharge-amazon-s3-with-amazon-elasticache-for-redis/>

upvoted 1 times

trinh_le 1 year, 1 month ago

Selected Answer: A

static content -> S3

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #895

Topic 1

A company is implementing a shared storage solution for a media application that the company hosts on AWS. The company needs the ability to use SMB clients to access stored data.

Which solution will meet these requirements with the LEAST administrative overhead?

- A. Create an AWS Storage Gateway Volume Gateway. Create a file share that uses the required client protocol. Connect the application server to the file share.
- B. Create an AWS Storage Gateway Tape Gateway. Configure tapes to use Amazon S3. Connect the application server to the Tape Gateway.
- C. Create an Amazon EC2 Windows instance. Install and configure a Windows file share role on the instance. Connect the application server to the file share.
- D. Create an Amazon FSx for Windows File Server file system. Connect the application server to the file system. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

LeonSauveterre 5 months, 1 week ago

Selected Answer: D

FSx for Windows File Server for sure. About other options:

A - Volume Gateway doesn't directly support SMB. SMB is more commonly supported in File Gateway or FSx for Windows.

B - Tape Gateway is for backup, not for file sharing via SMB.

C - This is actually a valid alternative if you're looking for more control, but it involves manual setup, maintenance, patches, backups and of course scaling, which increases administrative overhead in a big way.

upvoted 2 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: D

D. Create an Amazon FSx for Windows File Server file system. Connect the application server to the file system.

upvoted 2 times

Scheldon 1 year ago

Selected Answer: D

Answer D

<https://aws.amazon.com/fsx/windows/>
upvoted 4 times

trinh_le 1 year, 1 month ago

Selected Answer: D

SMB protocol -> FSx windows
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #896

Topic 1

A company is designing its production application's disaster recovery (DR) strategy. The application is backed by a MySQL database on an Amazon Aurora cluster in the us-east-1 Region. The company has chosen the us-west-1 Region as its DR Region.

The company's target recovery point objective (RPO) is 5 minutes and the target recovery time objective (RTO) is 20 minutes. The company wants to minimize configuration changes.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Create an Aurora read replica in us-west-1 similar in size to the production application's Aurora MySQL cluster writer instance.
- B. Convert the Aurora cluster to an Aurora global database. Configure managed failover. **Most Voted**
- C. Create a new Aurora cluster in us-west-1 that has Cross-Region Replication.
- D. Create a new Aurora cluster in us-west-1. Use AWS Database Migration Service (AWS DMS) to sync both clusters.

Correct Answer: B

Community vote distribution

B (100%)

Comments

LeonSauveterre 5 months, 1 week ago

Selected Answer: B

- A - Read replicas can't automatically provide failover in the event of a disaster.
- B - Aurora Global Databases are designed for cross-region replication and are synchronous, which means they can meet a low RPO (within seconds). Then failover is automated, meaning that if the primary region fails, traffic can be redirected to the secondary region automatically.
- C - Aurora does not automatically fail over with cross-region replication. You need to promote the replica manually, which could take longer than 20 minutes. The 5-minute RPO might not be met either.
- D - Well, just too much work to do manually. Let's simply rule it out.

upvoted 3 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: B

B. Convert the Aurora cluster to an Aurora global database. Configure managed failover.
upvoted 1 times

EdricHoang 11 months, 1 week ago

Selected Answer: B

I go for B. However, C is also a good option except manual failover intervention
upvoted 2 times

Scheldon 1 year ago

Selected Answer: B

Answer:B

<https://aws.amazon.com/rds/aurora/global-database/>

Cross-Region disaster recovery

If your primary Region suffers a performance degradation or outage, you can promote one of the secondary Regions to take read/write responsibilities. An Aurora cluster can recover in less than 1 minute, even in the event of a complete Regional outage. This provides your application with an effective recovery point objective (RPO) of 1 second and a recovery time objective (RTO) of less than 1 minute, providing a strong foundation for a global business continuity plan.

upvoted 3 times

sandordini 1 year, 1 month ago

Selected Answer: B

Aurora Global Database: allowing a single Amazon Aurora database to span multiple AWS Regions. It replicates your data with no impact on database performance, enables fast local reads with low latency in each Region, and provides disaster recovery from Region-wide outages.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #897

Topic 1

A company runs a critical data analysis job each week before the first day of the work week. The job requires at least 1 hour to complete the analysis. The job is stateful and cannot tolerate interruptions. The company needs a solution to run the job on AWS.

Which solution will meet these requirements?

- A. Create a container for the job. Schedule the job to run as an AWS Fargate task on an Amazon Elastic Container Service (Amazon ECS) cluster by using Amazon EventBridge Scheduler. **Most Voted**
- B. Configure the job to run in an AWS Lambda function. Create a scheduled rule in Amazon EventBridge to invoke the Lambda function.
- C. Configure an Auto Scaling group of Amazon EC2 Spot Instances that run Amazon Linux. Configure a crontab entry on the instances to run the analysis.
- D. Configure an AWS DataSync task to run the job. Configure a cron expression to run the task on a schedule.

Correct Answer: A

Community vote distribution

A (100%)

Comments

Hkayne **Highly Voted** 1 year, 1 month ago

Selected Answer: A

B is wrong because the job takes 1 hour and the lambda maximum execution time is 15 minutes.

C is wrong can't use spot instances because the job can not tolerate interruptions.

D is wrong too because DataSync is not designed to launch jobs.

Correct answer is A

upvoted 6 times

mk168898 **Most Recent** 7 months, 2 weeks ago

B is not because lambda only max 15mins but job requires 1 hour

C is not because cannot tolerate interruptions, spot instance might be gone anytime.

D datasync doesn't seem correct

A is most correct

upvoted 1 times

744fdad 10 months ago

I know datasync and lambda is event based, there are interruptions. C, doesnt address the scheduling requirement. Has to be A
upvoted 1 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: A

A. Create a container for the job. Schedule the job to run as an AWS Fargate task on an Amazon Elastic Container Service (Amazon ECS) cluster by using Amazon EventBridge Scheduler.

upvoted 1 times

Scheldon 1 year ago

Selected Answer: A

Answer: A

Fargate is compatibilie with ECS and is allowing for log running tasks

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #898

Topic 1

A company runs workloads in the AWS Cloud. The company wants to centrally collect security data to assess security across the entire company and to improve workload protection.

Which solution will meet these requirements with the LEAST development effort?

- A. Configure a data lake in AWS Lake Formation. Use AWS Glue crawlers to ingest the security data into the data lake.
- B. Configure an AWS Lambda function to collect the security data in .csv format. Upload the data to an Amazon S3 bucket.
- C. Configure a data lake in Amazon Security Lake to collect the security data. Upload the data to an Amazon S3 bucket.
- D. Configure an AWS Database Migration Service (AWS DMS) replication instance to load the security data into an Amazon RDS cluster.

Most Voted

Correct Answer: C

Community vote distribution

C (100%)

Comments

sandordini Highly Voted 1 year, 1 month ago

Selected Answer: C

A, B, D are senseless +

Amazon Security Lake automatically centralizes security data from AWS environments, you can get a more complete understanding of your security data across your entire organization. You can also improve the protection.

upvoted 7 times

LeonSauveterre Most Recent 5 months, 1 week ago

Selected Answer: C

Let's just focus on Security Lake.

1. Purpose-built: Amazon Security Lake is specifically designed to collect and centralize security data with minimal configuration and development overhead.
2. It integrates easily with AWS security services, providing a streamlined solution with less manual configuration compared to using AWS Glue or Lambda.
3. This is also designed to scale automatically, making it the most efficient and effective solution for the needs mentioned in the question.

upvoted 1 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: C

C. Configure a data lake in Amazon Security Lake to collect the security data. Upload the data to an Amazon S3 bucket.

upvoted 1 times

Scheldon 1 year ago

Selected Answer: C

Answer C

<https://aws.amazon.com/security-lake/>

Amazon Security Lake automatically centralizes security data from AWS environments, SaaS providers, on premises, and cloud sources into a purpose-built data lake stored in your account.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #899

Topic 1

A company is migrating five on-premises applications to VPCs in the AWS Cloud. Each application is currently deployed in isolated virtual networks on premises and should be deployed similarly in the AWS Cloud. The applications need to reach a shared services VPC. All the applications must be able to communicate with each other.

If the migration is successful, the company will repeat the migration process for more than 100 applications.

Which solution will meet these requirements with the LEAST administrative overhead?

- A. Deploy software VPN tunnels between the application VPCs and the shared services VPC. Add routes between the application VPCs in their subnets to the shared services VPC.
- B. Deploy VPC peering connections between the application VPCs and the shared services VPC. Add routes between the application VPCs in their subnets to the shared services VPC through the peering connection.
- C. Deploy an AWS Direct Connect connection between the application VPCs and the shared services VPC. Add routes from the application VPCs in their subnets to the shared services VPC and the applications VPCs. Add routes from the shared services VPC subnets to the applications VPCs.
- D. Deploy a transit gateway with associations between the transit gateway and the application VPCs and the shared services VPC. Add routes between the application VPCs in their subnets and the application VPCs to the shared services VPC through the transit gateway. **Most Voted**

Correct Answer: D*Community vote distribution*

D (93%)

B (7%)

Comments**Scheldon** **Highly Voted** 1 year ago**Selected Answer: D**

Answer: D

<https://aws.amazon.com/transit-gateway/>

Looks like the best solution would be transit gateway. It will allow for inter-VPC communication for all 5 applications/VPC, reach shared resource/VPC and in the future it will be easy to allow for inter-communication between even 100 VPCs

(applications)
upvoted 5 times

LeonSauveterre Most Recent 5 months, 1 week ago

Selected Answer: D

A - VPN Tunnels: High administrative overhead. Requires individual tunnel setup and routing configuration for each VPC, making scaling to 100+ VPCs inefficient.
B - VPC Peering: Moderate overhead. Requires manual peering for each pair of VPCs and routing configurations, which becomes harder to manage as more VPCs are added.
C - Direct Connect: High cost and complexity. Direct Connect is typically used for hybrid cloud setups or data center connectivity but is costly and harder to manage at large scale. Not suitable for connecting numerous VPCs in the AWS Cloud.
D - Transit Gateway: Lowest overhead. Acts as a centralized hub for managing inter-VPC communication and simplifies the process of routing between VPCs. It is highly scalable and reduces the complexity of adding new VPCs, and optimal for managing large-scale environments.
upvoted 1 times

LeonSauveterre 5 months, 1 week ago

I copied this from somewhere else for your references.
upvoted 1 times

mk168898 7 months, 2 weeks ago

each application needs to be in their own VPC and can communicate with each other => transit gateway
upvoted 2 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: D

D. Deploy a transit gateway with associations between the transit gateway and the application VPCs and the shared services VPC. Add routes between the application VPCs in their subnets and the application VPCs to the shared services VPC through the transit gateway.
upvoted 2 times

emakid 11 months, 2 weeks ago

Selected Answer: D

AWS Transit Gateway:

Centralized Connectivity: AWS Transit Gateway provides a hub-and-spoke model for connecting multiple VPCs, simplifying network management by providing a single point of connectivity for all VPCs.
Scalability: It is designed to handle many VPCs, making it suitable for scaling beyond the initial five applications to more than 100 applications.
Reduced Administrative Overhead: Managing VPC peering connections or VPN tunnels for each pair of VPCs would become complex and difficult to manage at scale. Transit Gateway simplifies this by providing centralized routing and connectivity.
upvoted 3 times

DanielWuTRT 11 months, 2 weeks ago

Selected Answer: D

the LEAST administrative overhead = transit gateway
upvoted 2 times

Obdf3af 1 year ago

D. <https://docs.aws.amazon.com/whitepapers/latest/building-scalable-secure-multi-vpc-network-infrastructure/transit-gateway.html>
upvoted 3 times

Hkayne 1 year, 1 month ago

Selected Answer: B

Correct answer is B
upvoted 1 times

dragosky 5 months, 3 weeks ago

not possible, vpc peerings can have up to 125 connections and the request is for 100 apps, meaning 100/100=1 1/2

not possible, vpc peering can have up to 125 connections and the request is for 100 apps, meaning $100 \times (100-1)/2$ connections...

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #900

Topic 1

A company wants to use Amazon Elastic Container Service (Amazon ECS) to run its on-premises application in a hybrid environment. The application currently runs on containers on premises.

The company needs a single container solution that can scale in an on-premises, hybrid, or cloud environment. The company must run new application containers in the AWS Cloud and must use a load balancer for HTTP traffic.

Which combination of actions will meet these requirements? (Choose two.)

- A. Set up an ECS cluster that uses the AWS Fargate launch type for the cloud application containers. Use an Amazon ECS Anywhere external launch type for the on-premises application containers. **Most Voted**
- B. Set up an Application Load Balancer for cloud ECS services. **Most Voted**
- C. Set up a Network Load Balancer for cloud ECS services.
- D. Set up an ECS cluster that uses the AWS Fargate launch type. Use Fargate for the cloud application containers and the on-premises application containers.
- E. Set up an ECS cluster that uses the Amazon EC2 launch type for the cloud application containers. Use Amazon ECS Anywhere with an AWS Fargate launch type for the on-premises application containers.

Correct Answer: AB

Community vote distribution

AB (100%)

Comments

Scheldon **Highly Voted** 1 year ago

Answer: AB

We need to load-balance HTTP traffic hence Application Load Balancer is needed. Because Customer want to use container solution we need to use ECS with Fargate which will lunch cloud applications. To run on-premises applications in containers we need to use ECS Anyware.

https://docs.aws.amazon.com/AmazonECS/latest/developerguide/AWS_Fargate.html

<https://docs.aws.amazon.com/AmazonECS/latest/developerguide/Welcome.html>

upvoted 5 times

Scheldon 1 year ago

Shreenu 1 year ago

Amazon ECS services on AWS Fargate support the Application Load Balancer and Network Load Balancer load balancer types. Application Load Balancers are used to route HTTP/HTTPS (or layer 7) traffic. Serverless (AWS Fargate (Fargate)) in the AWS cloud Fargate is a serverless, pay-as-you-go compute engine. With Fargate you don't need to manage servers, handle capacity planning, or isolate container workloads for security. On-premises virtual machines (VM) or servers Amazon ECS Anywhere provides support for registering an external instance such as an on-premises server or virtual machine (VM), to your Amazon ECS cluster.

upvoted 2 times

LeonSauveterre **Most Recent** 5 months, 1 week ago

Selected Answer: AB

E is incorrect because the combination of EC2 for cloud containers and Fargate for on-premises containers is unnecessary and would complicate the solution unnecessarily. Using ECS Anywhere with Fargate for both environments is a simpler and more efficient approach (option A).

Also, HTTP => ALB; TCP/UDP => NLB.

upvoted 1 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: AB

AB.the combination of setting up an ECS cluster with Fargate for the cloud and ECS Anywhere for on-premises, along with using an Application Load Balancer, provides a scalable, hybrid, and cloud-native solution with minimal operational overhead.

upvoted 2 times

emakid 11 months, 2 weeks ago

Selected Answer: AB

the combination of setting up an ECS cluster with Fargate for the cloud and ECS Anywhere for on-premises, along with using an Application Load Balancer, provides a scalable, hybrid, and cloud-native solution with minimal operational overhead.

upvoted 3 times

KennethNg923 11 months, 3 weeks ago

Selected Answer: AB

Have to use fargate, and since it has on-premises application so it need to use Amazon ECS Anywhere

upvoted 3 times

Obdf3af 1 year ago

BD.

<https://aws.amazon.com/blogs/aws/getting-started-with-amazon-ecs-anywhere-now-generally-available/>

upvoted 1 times

Hkayne 1 year, 1 month ago

Selected Answer: AB

AB is the correct answer,must lunch the cluster as external lunch

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #901

Topic 1

A company is migrating its workloads to AWS. The company has sensitive and critical data in on-premises relational databases that run on SQL Server instances.

The company wants to use the AWS Cloud to increase security and reduce operational overhead for the databases.

Which solution will meet these requirements?

- A. Migrate the databases to Amazon EC2 instances. Use an AWS Key Management Service (AWS KMS) AWS managed key for encryption.
- B. Migrate the databases to a Multi-AZ Amazon RDS for SQL Server DB instance. Use an AWS Key Management Service (AWS KMS) AWS managed key for encryption. **Most Voted**
- C. Migrate the data to an Amazon S3 bucket. Use Amazon Macie to ensure data security.
- D. Migrate the databases to an Amazon DynamoDB table. Use Amazon CloudWatch Logs to ensure data security.

Correct Answer: B

Community vote distribution

B (100%)

Comments

trinh_le **Highly Voted** 1 year, 1 month ago

Selected Answer: B

Migrate the databases to a Multi-AZ Amazon RDS for SQL Server DB instance. Use an AWS Key Management Service (AWS KMS) AWS managed key for encryption.

upvoted 6 times

mk168898 **Most Recent** 6 months, 4 weeks ago

looking to migrate SQL server isntance.

D is wrong because dynamoDB is noSQL

upvoted 2 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: B

B. Migrate the databases to a Multi-AZ Amazon RDS for SQL Server DB instance. Use an AWS Key Management Service (AWS KMS) AWS managed key for encryption.

upvoted 2 times

Seb888 10 months, 4 weeks ago

Selected Answer: B

B. Migrate the databases to a Multi-AZ Amazon RDS for SQL Server DB instance. Use an AWS Key Management Service (AWS KMS) AWS managed key for encryption.

This option provides the best balance between increased security, reduced operational overhead, and maintaining the relational database functionalities that the company needs.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #902

Topic 1

A company wants to migrate an application to AWS. The company wants to increase the application's current availability. The company wants to use AWS WAF in the application's architecture.

Which solution will meet these requirements?

A. Create an Auto Scaling group that contains multiple Amazon EC2 instances that host the application across two Availability Zones. Configure an Application Load Balancer (ALB) and set the Auto Scaling group as the target. Connect a WAF to the ALB. **Most Voted**

B. Create a cluster placement group that contains multiple Amazon EC2 instances that hosts the application. Configure an Application Load Balancer and set the EC2 instances as the targets. Connect a WAF to the placement group.

C. Create two Amazon EC2 instances that host the application across two Availability Zones. Configure the EC2 instances as the targets of an Application Load Balancer (ALB). Connect a WAF to the ALB.

D. Create an Auto Scaling group that contains multiple Amazon EC2 instances that host the application across two Availability Zones. Configure an Application Load Balancer (ALB) and set the Auto Scaling group as the target. Connect a WAF to the Auto Scaling group.

Correct Answer: A

Community vote distribution

A (100%)

Comments

trinh_le **Highly Voted** 1 year, 1 month ago

Selected Answer: A

Not D because AWS WAF cannot be directly connected to an Auto Scaling Group, it should be associated with the ALB which managing the incoming web traffic
upvoted 6 times

Scheldon **Highly Voted** 12 months ago

Selected Answer: A

AnswerA

WAF > ALB > AutoScalingGroup(MultiAZ EC2 Instances)

...with an Auto Scaling group (minimum 2 EC2 instances),

Looks good

upvoted 5 times

mk168898 Most Recent 6 months, 4 weeks ago

WAF manage incoming traffic should be at ALB, because auto scaling group is behind ALB then it defeats purpose of WAF

upvoted 2 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: A

A. Create an Auto Scaling group that contains multiple Amazon EC2 instances that host the application across two Availability Zones. Configure an Application Load Balancer (ALB) and set the Auto Scaling group as the target. Connect a WAF to the ALB.

upvoted 2 times

sandordini 1 year, 1 month ago

A: EC2 - MultiAZ > ALB > WAF

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #903

Topic 1


A company manages a data lake in an Amazon S3 bucket that numerous applications access. The S3 bucket contains a unique prefix for each application. The company wants to restrict each application to its specific prefix and to have granular control of the objects under each prefix.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create dedicated S3 access points and access point policies for each application. **Most Voted**
- B. Create an S3 Batch Operations job to set the ACL permissions for each object in the S3 bucket.
- C. Replicate the objects in the S3 bucket to new S3 buckets for each application. Create replication rules by prefix.
- D. Replicate the objects in the S3 bucket to new S3 buckets for each application. Create dedicated S3 access points for each application.

Correct Answer: A

Community vote distribution



Comments

anirudhsharma **Highly Voted** 1 year ago

Answer A

By creating separate access points for each application, you can enforce access controls specific to their respective prefixes while minimizing administrative complexity. This approach provides a clean separation of permissions and reduces the risk of misconfigurations.

Options B, C, and D are not as efficient or straightforward:

Option B (S3 Batch Operations) involves setting ACL permissions for each object individually, which can be cumbersome and time-consuming.

Option C (replicating objects to new S3 buckets) introduces additional buckets and replication rules, increasing management overhead.

Option D (replicating objects and creating dedicated S3 access points) adds unnecessary complexity by combining replication and access point creation.

upvoted 6 times

AMEJack **Most Recent** 6 months, 2 weeks ago

Selected Answer: A

"unique prefix for each application" hence access points. Also, ACL is not recommended by AWS.
upvoted 2 times

muhammadahmer36 10 months, 3 weeks ago

Selected Answer: A

A. Create dedicated S3 access points and access point policies for each application.

upvoted 2 times

emakid 11 months, 2 weeks ago

Selected Answer: A

Explanation:

S3 Access Points: These provide a way to manage access to shared data sets in Amazon S3. Each access point has a unique hostname and a policy that is specific to the use case, allowing for granular control over access to data.

Access Point Policies: These policies can be tailored to restrict access to specific prefixes within an S3 bucket, ensuring that each application only has access to its designated prefix.

upvoted 3 times

Scheldon 1 year ago

Answer B

Taking into consideration that we have "numerous applications" (10,100,1000?) and we need meet requirements with the LEAST operational overhead I would go into authomatization of operations hence Batch Operations seems to be good choice.

<https://aws.amazon.com/blogs/storage/updating-amazon-s3-object-acls-at-scale-with-s3-batch-operations/>

upvoted 1 times

f07ed8f 1 year ago

However, answer B need to set ACL for each OBJECT inside each S3... and the Batch operation has to be executed where there another new object added to the S3. I am not sure the answer too :)

upvoted 2 times

Tomrr 1 year ago

Selected Answer: A

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/access-points-policies.html>

upvoted 4 times

trinh_le 1 year, 1 month ago

Selected Answer: B

Create an S3 Batch Operations job to set the ACL permissions for each object in the S3 bucket

upvoted 2 times

aditianand 1 year ago

I have 2 questions: Is a batch process there to set ACL permissions. Secondly, they are asking with least operational overhead. Isn't A CORRECT?

upvoted 4 times

Topic 1