

FlyingHawk 4 months, 2 weeks ago

custom endpoint for an Aurora cluster represents a set of DB instances that you choose. When you connect to the endpoint, Aurora performs connection balancing and chooses one of the instances in the group to handle the connection. You define which instances this endpoint refers to, and you decide what purpose the endpoint serves.

An Aurora DB cluster has no custom endpoints until you create one. You can create up to five custom endpoints for each provisioned Aurora cluster or Aurora Serverless v2 cluster. You can't use custom endpoints for Aurora Serverless v1 clusters.

The custom endpoint provides balanced database connections based on criteria other than the read-only or read/write capability of the DB instances.

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: A

A - You can assign the three replicas with different compute and memory specifications to the custom endpoint, ensuring that queries are distributed only among these replicas.

B - Expensive, cumbersome, and unnecessary. The clone has separate storage and compute resources.

C - You can't target just the three replicas you want.

D - The reader endpoint automatically distributes read queries across *ALL* Aurora Replicas in the cluster.

upvoted 2 times

JA2018 6 months, 1 week ago

Selected Answer: A

This is why I choose Option A:

Using a Custom endpoint: This allows you to specify a specific set of Aurora replicas be used for the reporting queries, enabling you to target the three replicas with different compute and memory specifications while distributing the workload across them.

upvoted 1 times

JA2018 6 months, 1 week ago

Why the other options are not ideal:

- #B. Create a three-node cluster clone and use the reader endpoint: While you can create a separate cluster, using the "reader endpoint" would distribute queries across all replicas in that cluster, not just the three you want to target with different specs.

- #C. Use any of the instance endpoints for the selected three nodes: This approach would require manual connection management to each specific replica, not providing automatic load balancing across the chosen three nodes.

- #D. Use the reader endpoint to automatically distribute the read-only workload: The standard "reader endpoint" distributes reads across all Aurora replicas in the cluster, not just the selected three with specific configurations.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #947

Topic 1

A company runs a Node.js function on a server in its on-premises data center. The data center stores data in a PostgreSQL database. The company stores the credentials in a connection string in an environment variable on the server. The company wants to migrate its application to AWS and to replace the Node.js application server with AWS Lambda. The company also wants to migrate to Amazon RDS for PostgreSQL and to ensure that the database credentials are securely managed.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Store the database credentials as a parameter in AWS Systems Manager Parameter Store. Configure Parameter Store to automatically rotate the secrets every 30 days. Update the Lambda function to retrieve the credentials from the parameter.
- B. Store the database credentials as a secret in AWS Secrets Manager. Configure Secrets Manager to automatically rotate the credentials every 30 days. Update the Lambda function to retrieve the credentials from the secret. **Most Voted**
- C. Store the database credentials as an encrypted Lambda environment variable. Write a custom Lambda function to rotate the credentials. Schedule the Lambda function to run every 30 days.
- D. Store the database credentials as a key in AWS Key Management Service (AWS KMS). Configure automatic rotation for the key. Update the Lambda function to retrieve the credentials from the KMS key.

Correct Answer: B

Community vote distribution

B (100%)

Comments

[Removed] **Highly Voted** 9 months, 3 weeks ago

Selected Answer: B

Secrets Manager:

AWS Secrets Manager is specifically designed to store and manage sensitive information like database credentials. It provides built-in functionality for securely storing, retrieving, and automatically rotating credentials.

Automatic Rotation:

Secrets Manager can be configured to automatically rotate the database credentials at regular intervals (e.g., every 30 days). This reduces operational overhead by eliminating the need for manual credential rotation or custom rotation logic.

Integration with Lambda:

Secrets Manager can be configured to automatically rotate the database credentials at regular intervals (e.g., every 30 days). This reduces operational overhead by eliminating the need for manual credential rotation or custom rotation logic.

Lambda functions can easily retrieve credentials stored in Secrets Manager by calling the Secrets Manager API, which simplifies the application code and enhances security.

upvoted 5 times

LeonSauveterre Most Recent 5 months ago

Selected Answer: B

A - Parameter Store doesn't support automatic rotation.

C - Don't write a custom function when you can use off-the-rack ones.

D - AWS KMS is a service for managing encryption keys, not for managing secrets like database credentials. Also it doesn't provide direct support for managing or rotating database credentials.

upvoted 2 times

aragon_saa 9 months, 3 weeks ago

Selected Answer: B

Answer is B

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #948

Topic 1


A company wants to replicate existing and ongoing data changes from an on-premises Oracle database to Amazon RDS for Oracle. The amount of data to replicate varies throughout each day. The company wants to use AWS Database Migration Service (AWS DMS) for data replication. The solution must allocate only the capacity that the replication instance requires.

Which solution will meet these requirements?

- A. Configure the AWS DMS replication instance with a Multi-AZ deployment to provision instances across multiple Availability Zones.
- B. Create an AWS DMS Serverless replication task to analyze and replicate the data while provisioning the required capacity. **Most Voted**
- C. Use Amazon EC2 Auto Scaling to scale the size of the AWS DMS replication instance up or down based on the amount of data to replicate.
- D. Provision AWS DMS replication capacity by using Amazon Elastic Container Service (Amazon ECS) with an AWS Fargate launch type to analyze and replicate the data while provisioning the required capacity.

Correct Answer: B

Community vote distribution



Comments

[Removed] **Highly Voted** 9 months, 3 weeks ago

Selected Answer: B

B. Create an AWS DMS Serverless replication task to analyze and replicate the data while provisioning the required capacity.

Explanation:

AWS DMS Serverless is designed to automatically allocate and manage the necessary compute and memory resources based on the demand of the data replication workload. It scales capacity up or down according to the data replication requirements without manual intervention.

This approach ensures that the replication task uses only the required capacity at any given time, optimizing costs and resources, especially given that the amount of data to replicate varies throughout the day.

upvoted 9 times

JunsK1e **Highly Voted** 10 months, 1 week ago

Selected Answer: B

Selected Answer: B

Correct answer is B since the question need to allocate only the capacity that the replication instance requires
upvoted 5 times

EllenLiu **Most Recent** 5 months, 2 weeks ago

Selected Answer: B

https://docs.aws.amazon.com/dms/latest/userguide/CHAP_Serverless.html
upvoted 1 times

pujithacg8 10 months ago

correct answer could be "B"
upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: A

Answer is A
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #949

Topic 1

A company has a multi-tier web application. The application's internal service components are deployed on Amazon EC2 instances. The internal service components need to access third-party software as a service (SaaS) APIs that are hosted on AWS.

The company needs to provide secure and private connectivity from the application's internal services to the third-party SaaS application. The company needs to ensure that there is minimal public internet exposure.

Which solution will meet these requirements?

- A. Implement an AWS Site-to-Site VPN to establish a secure connection with the third-party SaaS provider.
- B. Deploy AWS Transit Gateway to manage and route traffic between the application's VPC and the third-party SaaS provider.
- C. Configure AWS PrivateLink to allow only outbound traffic from the VPC without enabling the third-party SaaS provider to establish.
- D. Use AWS PrivateLink to create a private connection between the application's VPC and the third-party SaaS provider.

Most Voted**Correct Answer: D***Community vote distribution*

D (100%)

Comments**LeonSauveterre** 5 months ago**Selected Answer: D**

C - PrivateLink exists because of its main benefit of bidirectional communication for secure access without public internet exposure. By doing this, you're restricting the SaaS provider's ability to establish sessions. Option C is misleading and just plain wrong. Option D alone is enough.

upvoted 4 times

c3c7e62 1 month ago

D is correct, just want to clarify, PrivateLink setup is unidirectional.

"AWS PrivateLink — Use AWS PrivateLink when you have a client/server set up where you want to allow one or more consumer VPCs unidirectional access to a specific service or set of instances in the service provider VPC or certain AWS services. Only the clients with access in the consumer VPC can initiate a connection to the service in the service provider VPC or AWS service. "

That means -

Inbound to service: Yes (from client to the endpoint service).

Outbound from service to client: No, unless the client explicitly exposes something .

refer to this -

<https://docs.aws.amazon.com/whitepapers/latest/building-scalable-secure-multi-vpc-network-infrastructure/aws-privatelink.html>

upvoted 1 times

spoved 8 months, 1 week ago

Selected Answer: D

<https://docs.aws.amazon.com/prescriptive-guidance/latest/integrate-third-party-services/architecture-1.html>

It is limited to only TCP traffic and unidirectional communication. The third-party workloads cannot initiate communication back to your account.

upvoted 4 times

56ce46c 8 months, 3 weeks ago

I think C is corret

2. Restrict Inbound Traffic via Security Groups:

To prevent the third-party SaaS provider from establishing inbound connections to your VPC, use Security Groups attached to the VPC Endpoint Interface.

Outbound Traffic Allowed: Ensure that your security groups allow outbound traffic to the SaaS provider's IP ranges or endpoints.

Restrict Inbound Traffic: You should block all inbound traffic on the VPC Endpoint Interface by configuring the security group rules. For example:

Inbound Rules: Block all traffic (or leave it empty).

Outbound Rules: Allow outbound connections to the IP addresses or ports specified by the SaaS provider.

upvoted 1 times

komorebi 10 months, 1 week ago

Selected Answer: D

Answer is D

upvoted 3 times

JunsK1e 10 months, 1 week ago

Selected Answer: D

D is correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #950

Topic 1


A solutions architect needs to connect a company's corporate network to its VPC to allow on-premises access to its AWS resources. The solution must provide encryption of all traffic between the corporate network and the VPC at the network layer and the session layer. The solution also must provide security controls to prevent unrestricted access between AWS and the on-premises systems.

Which solution meets these requirements?

- A. Configure AWS Direct Connect to connect to the VPC. Configure the VPC route tables to allow and deny traffic between AWS and on premises as required.
- B. Create an IAM policy to allow access to the AWS Management Console only from a defined set of corporate IP addresses. Restrict user access based on job responsibility by using an IAM policy and roles.
- C. Configure AWS Site-to-Site VPN to connect to the VPC. Configure route table entries to direct traffic from on premises to the VPC. Configure instance security groups and network ACLs to allow only required traffic from on premises. **Most Voted**
- D. Configure AWS Transit Gateway to connect to the VPC. Configure route table entries to direct traffic from on premises to the VPC. Configure instance security groups and network ACLs to allow only required traffic from on premises.

Correct Answer: C

Community vote distribution



Comments

Abbas_Abi_AWS **Highly Voted** 10 months ago

Selected Answer: C

AWS Direct Connect does not provide encryption by itself; it is often used in conjunction with VPN for encrypted traffic. Direct Connect primarily offers a dedicated connection and does not inherently satisfy the encryption requirement.

upvoted 8 times

LeonSauveterre **Most Recent** 5 months ago

Selected Answer: C

A - Direct Connect provides private connectivity but does not inherently include encryption.

B - IAM policies is irrelevant.

C - AWS Site-to-Site VPN uses IPsec (network layer) and SSL/TLS (session layer) to encrypt all traffic between the on-premises

C - Configure Site-to-Site VPN to connect to the VPC. Configure route table entries to direct traffic from on premises network and the AWS VPC.

D - Just like option A.

upvoted 2 times

Anyio 5 months, 1 week ago

Selected Answer: C

The correct answer is C. Configure AWS Site-to-Site VPN to connect to the VPC. Configure route table entries to direct traffic from on premises to the VPC. Configure instance security groups and network ACLs to allow only required traffic from on premises.

Explanation:

Option C: Correct. AWS Site-to-Site VPN offers encrypted network connections over the internet, providing encryption of all traffic at the network layer (IPsec) between the on-premises network and the VPC. Moreover, by using route tables, security groups, and network ACLs, you can carefully control the flow of traffic and restrict access, thereby meeting the requirement of preventing unrestricted access.

upvoted 2 times

blehbleh 8 months ago

Selected Answer: C

This is C, but not for all the reasons everyone is posting. D, also encrypts traffic and works at the network layer and also has security controls to prevent unrestricted access between AWS and on-premises systems.

So, if you thought D like I did initially you were very close. The reason it is C, is because C works at both the network and session layer while doing all the other requirements as well. Where as D only works at the network layer.

Happy studying!

upvoted 4 times

[Removed] 9 months, 3 weeks ago

Selected Answer: C

C is correct

upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: D

Answer is D

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Inter-Region gateway peering uses the same network infrastructure as VPC peering. Therefore traffic is encrypted using AES-256 encryption at the virtual network layer as it travels between Regions. Traffic is also encrypted using AES-256 encryption at the physical layer when it traverses network links that are outside of the physical control of AWS. As a result, traffic is double encrypted on network links outside the physical control of AWS. Within the same Region, traffic is encrypted at the physical layer only when it traverses network links that are outside of the physical control of AWS.

upvoted 1 times

JunsK1e 10 months, 1 week ago

Selected Answer: C

C is correct question needs to access between on prem and AWS

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #951

Topic 1

A company has a custom application with embedded credentials that retrieves information from a database in an Amazon RDS for MySQL DB cluster. The company needs to make the application more secure with minimal programming effort. The company has created credentials on the RDS for MySQL database for the application user.

Which solution will meet these requirements?

- A. Store the credentials in AWS Key Management Service (AWS KMS). Create keys in AWS KMS. Configure the application to load the database credentials from AWS KMS. Enable automatic key rotation
- B. Store the credentials in encrypted local storage. Configure the application to load the database credentials from the local storage. Set up a credentials rotation schedule by creating a cron job.
- C. Store the credentials in AWS Secrets Manager. Configure the application to load the database credentials from Secrets Manager. Set up a credentials rotation schedule by creating an AWS Lambda function for Secrets Manager. **Most Voted**
- D. Store the credentials in AWS Systems Manager Parameter Store. Configure the application to load the database credentials from Parameter Store. Set up a credentials rotation schedule in the RDS for MySQL database by using Parameter Store.

Correct Answer: C

Community vote distribution

C (100%)

Comments

[Removed] **Highly Voted** 9 months, 3 weeks ago

Selected Answer: C

C
 Explanation:
 AWS Secrets Manager is designed specifically for managing and automatically rotating credentials, including database credentials, API keys, and other secrets. It provides a secure and centralized place to store credentials and allows applications to retrieve them securely without hardcoding them in the application.
 Secrets Manager also offers built-in support for automatic rotation of credentials using Lambda functions, which reduces the manual effort needed for rotation and enhances security.
 This approach requires minimal programming effort because the application only needs to be configured to retrieve the credentials from Secrets Manager instead of being embedded within the application code.
 upvoted 5 times

LeonSauveterre Most Recent 5 months ago

Selected Answer: C

- A - AWS KMS is used for encryption and key management, not for storing and rotating credentials themselves.
- B - Why manually set up a cron job?
- C - Securely stores and manages sensitive information such as credentials, and provides native support for automatic rotation.
- D - Parameter Store doesn't provide automatic rotation.

upvoted 2 times

Anyio 5 months, 1 week ago

Selected Answer: C

The correct answer is C. Store the credentials in AWS Secrets Manager. Configure the application to load the database credentials from Secrets Manager. Set up a credentials rotation schedule by creating an AWS Lambda function for Secrets Manager.

Explanation:

Option C: Correct. AWS Secrets Manager is designed to securely store, manage, and retrieve database credentials. It supports automatic credentials rotation for Amazon RDS databases with minimal programming effort. Secrets Manager can automatically rotate the Amazon RDS database credentials using a built-in Lambda function, providing an integrated and secure solution for handling credentials.

upvoted 1 times

Omshanti 8 months, 1 week ago

Selected Answer: C

AWS Secret manager securely stores data base user id and passwords

upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: C

Answer is C

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #952

Topic 1


A company wants to move its application to a serverless solution. The serverless solution needs to analyze existing data and new data by using SQL. The company stores the data in an Amazon S3 bucket. The data must be encrypted at rest and replicated to a different AWS Region.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create a new S3 bucket that uses server-side encryption with AWS KMS multi-Region keys (SSE-KMS). Configure Cross-Region Replication (CRR). Load the data into the new S3 bucket. Use Amazon Athena to query the data. **Most Voted**
- B. Create a new S3 bucket that uses server-side encryption with Amazon S3 managed keys (SSE-S3). Configure Cross-Region Replication (CRR). Load the data into the new S3 bucket. Use Amazon RDS to query the data.
- C. Configure Cross-Region Replication (CRR) on the existing S3 bucket. Use server-side encryption with Amazon S3 managed keys (SSE-S3). Use Amazon Athena to query the data.
- D. Configure S3 Cross-Region Replication (CRR) on the existing S3 bucket. Use server-side encryption with AWS KMS multi-Region keys (SSE-KMS). Use Amazon RDS to query the data.

Correct Answer: A

Community vote distribution



Comments

dhewa **Highly Voted** 9 months, 3 weeks ago

Selected Answer: A

A wins because it gives us encryption with AWS KMS multi-Region keys
upvoted 7 times

LeonSauveterre **Highly Voted** 5 months ago

Selected Answer: C

A - A new bucket + KMS multi-Region keys = Too much operational oversight.
 B - RDS is not a serverless solution
 C - By using the existing S3 bucket, you eliminate the need to create a new bucket and load data into it.
 D - Athena supports querying using SQL, so that rules out RDS.
 upvoted 5 times

El189 Most Recent 2 months ago

Selected Answer: A

A is correct.

everyone. voting for c i think isnt paying attention to the part where you need the data replicated to a different region. how are you meant to replicate it without a new bucket?

everyone saying use the existing bucket.... urm what? how can you replicate into the same bucket thats in the same region and expect it to be in a different region?

clearly A.

upvoted 1 times

tch 2 months, 3 weeks ago

Selected Answer: C

we don't need to create a policy specifically for Amazon S3-managed keys (SSE-S3) to encrypt objects in your S3 bucket... should be more easy and LEAST operational overhead

upvoted 1 times

yangbo 3 months ago

Selected Answer: C

The default encryption for S3 is SSE-S3, and cross-region replication can be enabled normally.

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: A

Amazon S3 managed keys is region specific, for CRR, we must use KMS multi-region keys.
<https://docs.aws.amazon.com/kms/latest/developerguide/multi-region-keys-overview.html>

upvoted 5 times

FlyingHawk 4 months, 2 weeks ago

I checked AI, ChatGPT, Gemini , Claudi AI and DeepSeek think A is correct, but meta AI thinks C is correct. SSE-S3 is the default encryption, it will be least operational effort for sure. but with SSE-S3 (server-side encryption with S3-managed keys), AWS manages the encryption keys, and each region will use its own region-specific encryption key.

There is no shared key between the source and destination buckets. AWS handles encryption and decryption seamlessly in each region. This approach is operationally simple but lacks control and consistency for encryption keys across regions.

upvoted 2 times

trinh_le 5 months ago

Selected Answer: C

It's not require highly sensitive data or complexity gain permission. So it should use sse-s3 is suitable

upvoted 1 times

Anyio 5 months, 2 weeks ago

Selected Answer: C

The correct answer is C.

Explanation:

Option A: Incorrect. Although using AWS KMS multi-Region keys (SSE-KMS) and Amazon Athena to query the data meet the security and SQL querying requirements, creating a new S3 bucket and handling data migration increases operational overhead unnecessarily if the data already exists.

Option C: Correct. Configuring Cross-Region Replication (CRR) on the existing S3 bucket makes efficient use of existing infrastructure and leverages server-side encryption with Amazon S3 managed keys (SSE-S3) to ensure data encryption at rest with lower operational complexity and costs compared to using KMS keys. Using Amazon Athena allows querying the data directly in S3, offering serverless and flexible SQL querying capabilities with minimal setup and operational overhead.

upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: C

Selected Answer: C

I will chose Option C for the following reasons:

#1: Least operational overhead: Choosing "SSE-S3" over "SSE-KMS" minimizes operational overhead as it automatically manages encryption keys within S3, eliminating the need for additional KMS key management.

#2: Existing S3 bucket: Reusing the existing bucket avoids the extra step of creating a new one and migrating data.

#3: Athena for querying: Athena is a serverless solution ideal for querying large datasets stored in S3, aligning with the requirement for a serverless architecture.

upvoted 2 times

JA2018 6 months, 1 week ago

For Option A: While using KMS multi-region keys provides more control, it adds extra management complexity compared to SSE-S3.

upvoted 2 times

Abdullah2004 9 months, 1 week ago

Selected Answer: A

A is correct

upvoted 3 times

komorebi 10 months, 1 week ago

Selected Answer: A

Answer is A

upvoted 4 times

JunsK1e 10 months, 1 week ago

Selected Answer: C

C is correct because it needs to replicate to different AWS region

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #953

Topic 1

A company has a web application that has thousands of users. The application uses 8-10 user-uploaded images to generate AI images. Users can download the generated AI images once every 6 hours. The company also has a premium user option that gives users the ability to download the generated AI images anytime.


The company uses the user-uploaded images to run AI model training twice a year. The company needs a storage solution to store the images.

Which storage solution meets these requirements MOST cost-effectively?

- A. Move uploaded images to Amazon S3 Glacier Deep Archive. Move premium user-generated AI images to S3 Standard. Move non-premium user-generated AI images to S3 Standard-Infrequent Access (S3 Standard-IA). **Most Voted**
- B. Move uploaded images to Amazon S3 Glacier Deep Archive. Move all generated AI images to S3 Glacier Flexible Retrieval.
- C. Move uploaded images to Amazon S3 One Zone-Infrequent Access (S3 One Zone-IA). Move premium user-generated AI images to S3 Standard. Move non-premium user-generated AI images to S3 Standard-Infrequent Access (S3 Standard-IA).
- D. Move uploaded images to Amazon S3 One Zone-Infrequent Access (S3 One Zone-IA). Move all generated AI images to S3 Glacier Flexible Retrieval.

Correct Answer: A

Community vote distribution



Comments

dhewa **Highly Voted** 9 months, 3 weeks ago

Selected Answer: A

The company uses the user-uploaded images to run AI model training twice a year. So for this Deep Archive will be necessary. upvoted 7 times

JA2018 6 months, 1 week ago

But moving uploaded images to Glacier Deep Archive is not ideal as it has a very high retrieval cost and long retrieval time, which is not suitable for the occasional access needed for AI model training.

upvoted 2 times

blehbleh Highly Voted 8 months ago

Selected Answer: A

This is A. We care about cost effectiveness. In regards to images being used twice a year "For images accessed only twice a year, S3 Glacier Deep Archive would be the more cost-effective option compared to S3 One Zone Infrequent Access, as it is designed for extremely infrequent access and offers the lowest storage cost within AWS S3 storage classes; while S3 One Zone Infrequent Access is cheaper than standard Infrequent Access, it still might be slightly more expensive for data accessed as rarely as twice a year." In regards to premium users keep them standard. Non premium users can be in the infrequent since they have 6 hrs.

Cost Effective!

upvoted 5 times

JA2018 6 months, 1 week ago

But moving uploaded images to Glacier Deep Archive is not ideal as it has a very high retrieval cost and long retrieval time, which is not suitable for the occasional access needed for AI model training.

upvoted 1 times

tch Most Recent 3 months ago

Selected Answer: A

S3 Glacier Deep Archive is Amazon S3's lowest-cost storage class and supports long-term retention and digital preservation for data that may be accessed once or twice in a year

B& D are out.....

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: C

A - Deep Archive is designed for long-term cold storage and can spend hours to retrieve each object. For AI model training, you'd hate it if you must spend hundreds/thousands of hours before you can actually start training it.

B - Premium users would be so mad if they had to wait for hours to get their generated images.

C - S3 One Zone-IA reduces costs, and it's great now that we don't need high durability across multiple availability zones.

Premium users will experience low latency to get their generated images.

D - Glacier tiers are not designed for frequent access. This is bad practice.

One more thing about option C, since training happens extremely infrequently and the images can be re-uploaded if necessary, One Zone IA's lower redundancy is an acceptable trade-off for cost savings.

upvoted 2 times

Cpso 6 months, 1 week ago

Selected Answer: A

I choose A. But it not best answer. for standard-ia you get data immediately but pay for request which not make sense for requirement. The best fit of non-premium user is Glacier Fleixble standard(1-5hour)

upvoted 3 times

JA2018 6 months, 1 week ago

But moving uploaded images to Glacier Deep Archive is not ideal as it has a very high retrieval cost and long retrieval time, which is not suitable for the occasional access needed for AI model training.

upvoted 1 times

Oghare 7 months, 4 weeks ago

this is A

S3 Glacier Deep Archive is the most cost-effective storage for long-term infrequently accessed data, ideal for the user-uploaded images used twice a year for AI training.

S3 Standard for premium user-generated AI images ensures fast access, which is needed for frequent downloads.

S3 Standard-IA for non-premium user-generated AI images is cost-effective for less frequent access, as it charges lower for storage but slightly more for retrieval, which fits the 6-hour download frequency.

Not C because, although S3 One Zone-IA is a lower-cost option as well it provides less durability because it stores data in only one Availability Zone. While it is cost-effective, it increases the risk of data loss for critical AI training data.

B and D is out

S3 Glacier Flexible Retrieval for all generated AI images would likely introduce unacceptable retrieval delays for premium users, as they require immediate access to download images

upvoted 3 times

JA2018 6 months, 1 week ago

But moving uploaded images to Glacier Deep Archive is not ideal as it has a very high retrieval cost and long retrieval time, which is not suitable for the occasional access needed for AI model training.

upvoted 1 times

kbgsqsgs 8 months, 1 week ago

Selected Answer: C

S3 Glacier Deep Archive does not meet the 6-hour download requirement because it takes time to access data. User can download even if you are not a premium user

upvoted 1 times

blehbleh 8 months ago

it says "Users can download the generated AI images once every 6 hours." The upload images are used for training only 2 times a year which meets the s3 glacier deep dive. The standard users are put in infrequent access and premium users are put in standard. I don't understand how people are messing this up.

upvoted 3 times

[Removed] 9 months, 3 weeks ago

Selected Answer: B

B is correct

Explanation:

S3 Glacier Deep Archive is the most cost-effective storage option for data that is rarely accessed. Since the user-uploaded images are only used twice a year for AI model training, storing them in Glacier Deep Archive is ideal for minimizing costs. The longer retrieval time (up to 12 hours) is acceptable given the infrequent access.

S3 Glacier Flexible Retrieval is suitable for storing the generated AI images because it balances cost and retrieval time. Regular users can download images every 6 hours, which Glacier Flexible Retrieval can accommodate with its flexible retrieval options (ranging from minutes to hours). This solution also works for premium users, who might need more frequent access. While S3 Standard or Standard-IA could be used, Glacier Flexible Retrieval offers significant cost savings while still meeting the access requirements.

upvoted 2 times

JA2018 6 months, 1 week ago

Option B is using Glacier Flexible Retrieval.

Storing all generated AI images in Glacier Flexible Retrieval is also not cost-effective as the premium users require faster access to their images.

upvoted 1 times

officedepotadmin 9 months, 3 weeks ago

Selected Answer: C

S3 One Zone-IA is a cost-effective storage option for images that are accessed infrequently but are still needed for AI model training twice a year. One Zone-IA stores data in a single Availability Zone, making it less expensive but still highly available within that zone. Premium users need frequent access to their AI-generated images so S3. Non-premium users access their AI-generated images less frequently (once every 6 hours) so S3 Standard-IA

upvoted 5 times

pujithacg8 10 months ago

A is correct as Glacier deep archive provides the lowest-cost storage class.

upvoted 1 times

JA2018 6 months, 1 week ago

But moving uploaded images to Glacier Deep Archive is not ideal as it has a very high retrieval cost and long retrieval time, which is not suitable for the occasional access needed for AI model training.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #954

Topic 1

A company is developing machine learning (ML) models on AWS. The company is developing the ML models as independent microservices. The microservices fetch approximately 1 GB of model data from Amazon S3 at startup and load the data into memory. Users access the ML models through an asynchronous API. Users can send a request or a batch of requests.

The company provides the ML models to hundreds of users. The usage patterns for the models are irregular. Some models are not used for days or weeks. Other models receive batches of thousands of requests at a time.


Which solution will meet these requirements?

- A. Direct the requests from the API to a Network Load Balancer (NLB). Deploy the ML models as AWS Lambda functions that the NLB will invoke. Use auto scaling to scale the Lambda functions based on the traffic that the NLB receives.
- B. Direct the requests from the API to an Application Load Balancer (ALB). Deploy the ML models as Amazon Elastic Container Service (Amazon ECS) services that the ALB will invoke. Use auto scaling to scale the ECS cluster instances based on the traffic that the ALB receives.
- C. Direct the requests from the API into an Amazon Simple Queue Service (Amazon SQS) queue. Deploy the ML models as AWS Lambda functions that SQS events will invoke. Use auto scaling to increase the number of vCPUs for the Lambda functions based on the size of the SQS queue.

D. Direct the requests from the API into an Amazon Simple Queue Service (Amazon SQS) queue. Deploy the ML models as Amazon Elastic Container Service (Amazon ECS) services that read from the queue. Use auto scaling for Amazon ECS to scale both the cluster capacity and number of the services based on the size of the SQS queue. **Most Voted**

Correct Answer: D

Community vote distribution

**Comments**

[Removed] **Highly Voted** 9 months, 3 weeks ago

Selected Answer: D

D is the answer:

SQS Queue: Directing API requests to SQS decouples the API from ML processing, efficiently handles high traffic, and ensures reliable request processing without overloading the ML models.

Amazon ECS Services: Running ML models on ECS provides effective management of containerized applications, ideal for handling ML workloads.

Auto Scaling: ECS auto scales based on SQS queue size, adjusting container and cluster capacity to match demand, ensuring efficient handling of varying workloads.

upvoted 5 times

FlyingHawk Most Recent 4 months, 3 weeks ago

Selected Answer: D

Given that the problem statement does not specify the processing time, Option D (ECS) is the safer choice because it does not have the same limitations as Lambda(cold start and 15 ,minutes execution). However, if the processing time is short and cold starts are acceptable, Option C (Lambda) could also be a cost-effective solution.

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: D

A - Lambda has a cold start latency. There are ways to optimize that (like the environment variable config which I actually did when I used Lambda once), but loading 1GB into memory is still too much.

B - Using ALB makes handling sudden bursts of traffic even more challenging.

C - Just like A if I'm correct.

D - Each ML model needs to load 1 GB of model data at startup ==> Containers in ECS are well-suited for this because they provide persistent execution environments that avoid the overhead of reloading the model data frequently.

upvoted 1 times

RamanAgarwal 5 months, 3 weeks ago

Selected Answer: C

Since AWS Lambda supports 1GB of memory and can be scaled seamlessly, the ans should be C. Question says a lot of APIs are not used frequently so keeping them in ECS will result in higher costs. Also these are async operations means the response is not time bound hence we can wait for lambda to startup and scale up based on the size of the SQS queue.

upvoted 2 times

kbgsqsgs 8 months, 1 week ago

Selected Answer: B

Why should I use SQS in option D? Wouldn't ALB be enough?

upvoted 2 times

Sergantus 6 months, 3 weeks ago

It's talking about accessing models through an asynchronous API, so decoupling is needed (SQS)

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #955

Topic 1

A company runs a web application on Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer (ALB). The application stores data in an Amazon Aurora MySQL DB cluster.

The company needs to create a disaster recovery (DR) solution. The acceptable recovery time for the DR solution is up to 30 minutes. The DR solution does not need to support customer usage when the primary infrastructure is healthy.

Which solution will meet these requirements?

A. Deploy the DR infrastructure in a second AWS Region with an ALB and an Auto Scaling group. Set the desired capacity and maximum capacity of the Auto Scaling group to a minimum value. Convert the Aurora MySQL DB cluster to an Aurora global database. Configure Amazon Route 53 for an active-passive failover with ALB endpoints. **Most Voted**


B. Deploy the DR infrastructure in a second AWS Region with an ALB. Update the Auto Scaling group to include EC2 instances from the second Region. Use Amazon Route 53 to configure active-active failover. Convert the Aurora MySQL DB cluster to an Aurora global database.

C. Back up the Aurora MySQL DB cluster data by using AWS Backup. Deploy the DR infrastructure in a second AWS Region with an ALB. Update the Auto Scaling group to include EC2 instances from the second Region. Use Amazon Route 53 to configure active-active failover. Create an Aurora MySQL DB cluster in the second Region. Restore the data from the backup.

D. Back up the infrastructure configuration by using AWS Backup. Use the backup to create the required infrastructure in a second AWS Region. Set the Auto Scaling group desired capacity to zero. Use Amazon Route 53 to configure active-passive failover. Convert the Aurora MySQL DB cluster to an Aurora global database.

Correct Answer: A

Community vote distribution



Comments

chwieobjom **Highly Voted** 9 months, 3 weeks ago

Selected Answer: A

RTO 30 minute, warm standby.

upvoted 12 times

upvoted 13 times

spoved Highly Voted 8 months, 1 week ago

Selected Answer: D

<https://docs.aws.amazon.com/whitepapers/latest/disaster-recovery-workloads-on-aws/disaster-recovery-options-in-the-cloud.html>

The DR solution does not need to support customer usage when the primary infrastructure is healthy. -> Pilot Light

upvoted 7 times

jingen11 7 months, 3 weeks ago

Backup and restore RTO is up to hours

upvoted 3 times

zdi561 Most Recent 3 months, 3 weeks ago

Selected Answer: D

A is not right because the max and min are set to minimum which is zero. and the global table use resource.

upvoted 1 times

Salilgen 4 months, 4 weeks ago

Selected Answer: A

The question doesn't ask about the most cost effective solution. A is the option that ensures the lowest recovery time

upvoted 3 times

LeonSauveterre 5 months ago

Selected Answer: A

Well just a thought, I hate it when each option looks like a valid one... I think I'll go with A because it seems to me other options are somewhat flawed compared to A.

B - "The DR solution does not need to support customer usage when the primary infrastructure is healthy." So active-active failover is not necessary.

C - Restoring a database from AWS Backup can take hours to complete, let alone the 30-minute RTO. If I remember it correctly, it can take up to 7 days?

D - Just like C.

One more thing about option D, the desired capacity parameter in an Auto Scaling group determines the number of EC2 instances that should be running. By setting it to 0, the group will not launch any EC2 instances during normal operations, so costs are reduced when idle. That's a plus actually.

upvoted 4 times

Anyio 5 months, 1 week ago

Selected Answer: A

The correct answer is A.

Explanation:

Option A: Correct. This solution sets up disaster recovery infrastructure in a secondary region with minimal running costs by keeping the Auto Scaling group's desired and maximum capacity low. By converting the Aurora MySQL DB cluster to an Aurora global database, the company can provide rapid cross-region replication and promote the standby database to primary if needed, minimizing downtime. Configuring active-passive failover using Route 53 ensures the DR system is only activated when necessary, aligning with the requirement of not supporting customer usage when the primary infrastructure is healthy.

upvoted 3 times

Liongeek 5 months, 1 week ago

Selected Answer: A

RTO 30 min has to be warm standby.

AWS Backup DOES NOT back up configuration, it backs up DATA

upvoted 5 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: D

pilot light, set auto scaling group desired capacity to 0

upvoted 1 times

Penjerla 6 months ago

Selected Answer: A

Restoring Database from AWS back up could take from a few min to several hours. So not D

upvoted 3 times

JA2018 6 months, 1 week ago

Selected Answer: A

I will choose Option A for the following options:

#1 - Active-passive failover: This aligns with the requirement of a DR solution that doesn't need to support customer usage during normal operations, meaning the secondary region only activates when a primary failure occurs.

2 - Minimal Auto Scaling group capacity: Setting the desired and maximum capacity to a minimum value ensures minimal cost when the DR infrastructure is not actively used.

3 - Aurora Global Database: Converting the Aurora cluster to a global database allows for fast failover to the secondary region, meeting the 30-minute recovery time objective.

#4 - Separate AWS Region: Deploying the DR infrastructure in a different region provides geographic redundancy, crucial for disaster recovery.

upvoted 3 times

Atdotcom 7 months ago

Selected Answer: D

D is pilot-light because desired min capacity =0 so ressources are live , services idle, RTO 10s.Ressources are provisionned and scale after even.

upvoted 2 times

elmyth 7 months, 3 weeks ago

Selected Answer: D

It can be A, if minimum value = 0.

It Can be D, if this approach is implemented before the disaster.

And "Set the Auto Scaling group desired capacity to zero" means that it supposed to be done before the disaster.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #956


Topic 1

A company is migrating its data processing application to the AWS Cloud. The application processes several short-lived batch jobs that cannot be disrupted. Data is generated after each batch job is completed. The data is accessed for 30 days and retained for 2 years.

The company wants to keep the cost of running the application in the AWS Cloud as low as possible.

Which solution will meet these requirements?

- A. Migrate the data processing application to Amazon EC2 Spot Instances. Store the data in Amazon S3 Standard. Move the data to Amazon S3 Glacier Instant. Retrieval after 30 days. Set an expiration to delete the data after 2 years.
- B. Migrate the data processing application to Amazon EC2 On-Demand Instances. Store the data in Amazon S3 Glacier Instant Retrieval. Move the data to S3 Glacier Deep Archive after 30 days. Set an expiration to delete the data after 2 years.
- C. Deploy Amazon EC2 Spot Instances to run the batch jobs. Store the data in Amazon S3 Standard. Move the data to Amazon S3 Glacier Flexible Retrieval after 30 days. Set an expiration to delete the data after 2 years.
- D. Deploy Amazon EC2 On-Demand Instances to run the batch jobs. Store the data in Amazon S3 Standard. Move the data to Amazon S3 Glacier Deep Archive after 30 days. Set an expiration to delete the data after 2 years. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****nebajp** **Highly Voted** 10 months ago**Selected Answer: D**

D is the correct answer
for 30 days - use Amazon S3 standard
2 years Retaining - Glacier Deep Archive
Can not be Disrupted - On-Demand Instances
upvoted 10 times

flaviobrf **Highly Voted** 10 months ago

Selected Answer: D

I understand that D is the right answer

upvoted 5 times

FlyingHawk **Most Recent** 4 months, 3 weeks ago

Selected Answer: D

1. The short-lived batch jobs cannot be disrupted -> on demand. A and C are out
2. The data is accessed for 30 days and retained for 2 years. means S3 Standard (30 days) -> S3 Glacier -> expire. S3 Glacier Instant Retrieval has a minimum storage for 90 days, so B is out.

upvoted 3 times

LeonSauveterre 5 months ago

Selected Answer: D

B - Amazon S3 Glacier Instant Retrieval is an archive storage class that delivers the lowest-cost storage for long-lived data that is rarely accessed and requires retrieval in milliseconds. So this is wrong because "The data is accessed for 30 days".

upvoted 2 times

Anyio 5 months, 1 week ago

Selected Answer: D

The correct answer is D. Deploy Amazon EC2 On-Demand Instances to run the batch jobs. Store the data in Amazon S3 Standard. Move the data to Amazon S3 Glacier Deep Archive after 30 days. Set an expiration to delete the data after 2 years.

Explanation:

Option D: Correct. Since the batch jobs cannot be disrupted, using Amazon EC2 On-Demand Instances ensures the application runs without interruption. Initially storing data in Amazon S3 Standard allows for access during the 30-day period post-processing. Moving the data to Amazon S3 Glacier Deep Archive after 30 days optimizes storage costs for long-term retention of up to 2 years, the lowest cost for infrequent access. Setting an expiration ensures the data is deleted after 2 years, controlling costs and data lifecycle management.

upvoted 2 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: D

I am always confused about the choice between s3 standard and s3 glacier instant...

short period of storage + frequently retrieve ==> s3 standard

long period of storage + infrequently retrieve ==> s3 glacier instant

s3 standard: \$0.023 per GB for storage / \$0.0004 get, select per 1000 request

s3 glacier instant: \$0.004 per GB for storage / \$0.01 get, select per 1000 request

<https://aws.amazon.com/s3/pricing/?nc=sn&loc=4>

upvoted 3 times

spoved 8 months, 1 week ago

Selected Answer: B

<https://aws.amazon.com/about-aws/whats-new/2021/11/amazon-s3-glacier-instant-retrieval-storage-class/>

- The easiest way to store data in S3 Glacier Instant Retrieval is to use the S3 PUT API to upload data directly, or use S3 Lifecycle to transition data from the S3 Standard and S3 Standard-IA storage classes.

- The company wants to keep the cost of running the application in the AWS Cloud as low as possible

=> B

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

S3 Glacier Instant Retrieval has a minimum storage duration of 90 days. If you move data to S3 Glacier Deep Archive before 90 days, you will be charged for the full 90 days of storage, even if the data is deleted or moved earlier. This makes it cost-ineffective for the company's requirement to move data after 30 days.

upvoted 1 times

[Removed] 9 months, 3 weeks ago

Selected Answer: D

D looks right

upvoted 4 times

pujithacg8 10 months ago

D is correct
upvoted 2 times

siheom 10 months, 1 week ago

Selected Answer: C

I VOTE C
upvoted 1 times

tch 3 months ago

A& C are out!
Answer is D
S3 Glacier Deep Archive is Amazon S3's lowest-cost storage class and supports long-term retention and digital preservation for data that may be accessed once or twice in a year.
upvoted 1 times

officedepotadmin 9 months, 3 weeks ago

you voted wrong
upvoted 1 times

SR0312 10 months, 1 week ago

Selected Answer: B

Job cannot be disrupted - On demand
upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #957

Topic 1

A company needs to design a hybrid network architecture. The company's workloads are currently stored in the AWS Cloud and in on-premises data centers. The workloads require single-digit latencies to communicate. The company uses an AWS Transit Gateway transit gateway to connect multiple VPCs.

Which combination of steps will meet these requirements MOST cost-effectively? (Choose two.)

A. Establish an AWS Site-to-Site VPN connection to each VPC.

B. Associate an AWS Direct Connect gateway with the transit gateway that is attached to the VPCs. **Most Voted**


C. Establish an AWS Site-to-Site VPN connection to an AWS Direct Connect gateway.

D. Establish an AWS Direct Connect connection. Create a transit virtual interface (VIF) to a Direct Connect gateway. **Most Voted**

E. Associate AWS Site-to-Site VPN connections with the transit gateway that is attached to the VPCs.

Correct Answer: BD

Community vote distribution



Comments

[Removed] **Highly Voted** 9 months, 3 weeks ago

Selected Answer: BD

BD sounds right
upvoted 5 times

LeonSauveterre **Most Recent** 5 months ago

Selected Answer: BD

A little memo for me. Hope it'll be helpful for you as well.

1. VPN Connection: For secure, encrypted traffic between on-premises and AWS over the public internet. ==> Relies on public internet (higher latency)
2. Direct Connect: For dedicated private network connectivity between on-premises and AWS (like data centers with large workloads). ==> Higher cost, longer setup time.
3. PrivateLink: For secure private connectivity to AWS services or third-party services (like Salesforce) within AWS. ==> Simple setup but limited to services configured with PrivateLink (not general connectivity).
4. VPC Peering: For direct connectivity between two AWS VPCs (like multi-account or region scenarios). ==> Does not support

4. VPC Peering: for direct connectivity between two AWS VPCs (like multi-account or region scenarios). --- Does not support transitive connections; not suitable for external services like Salesforce.

upvoted 3 times

MSAM16 5 months, 1 week ago

Selected Answer: BD

BD is correct

upvoted 1 times

jingen11 7 months, 3 weeks ago

Selected Answer: BE

Question asked about cheapest. D is not cheap

upvoted 1 times

Sergantus 6 months, 3 weeks ago

Single-digit latency is required

upvoted 3 times

muhammadahmer36 9 months, 4 weeks ago

Selected Answer: BD

BD is correct

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #958

Topic 1

A global ecommerce company runs its critical workloads on AWS. The workloads use an Amazon RDS for PostgreSQL DB instance that is configured for a Multi-AZ deployment.

Customers have reported application timeouts when the company undergoes database failovers. The company needs a resilient solution to reduce failover time.

Which solution will meet these requirements?

- A. Create an Amazon RDS Proxy. Assign the proxy to the DB instance. **Most Voted**
- B. Create a read replica for the DB instance. Move the read traffic to the read replica.
- C. Enable Performance Insights. Monitor the CPU load to identify the timeouts.
- D. Take regular automatic snapshots. Copy the automatic snapshots to multiple AWS Regions.

Correct Answer: A

Community vote distribution

A (100%)

Comments

[Removed] **Highly Voted** 9 months, 3 weeks ago

Selected Answer: A

A. Create an Amazon RDS Proxy. Assign the proxy to the DB instance.

Explanation:
Amazon RDS Proxy:

RDS Proxy is designed to manage connections to the database more efficiently. It can reduce the impact of failovers on the application by maintaining connections and transparently rerouting them to the standby instance during a failover event. By using RDS Proxy, the failover time is reduced because the proxy minimizes the disruption that occurs when the database fails over, thus reducing application timeouts.

upvoted 5 times

aragon_saa **Most Recent** 9 months, 3 weeks ago

Selected Answer: A

Answer is A
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #959

Topic 1

A company has multiple Amazon RDS DB instances that run in a development AWS account. All the instances have tags to identify them as development resources. The company needs the development DB instances to run on a schedule only during business hours.


Which solution will meet these requirements with the LEAST operational overhead?

- A. Create an Amazon CloudWatch alarm to identify RDS instances that need to be stopped. Create an AWS Lambda function to start and stop the RDS instances.
- B. Create an AWS Trusted Advisor report to identify RDS instances to be started and stopped. Create an AWS Lambda function to start and stop the RDS instances.
- C. Create AWS Systems Manager State Manager associations to start and stop the RDS instances.

D. Create an Amazon EventBridge rule that invokes AWS Lambda functions to start and stop the RDS instances. **Most Voted**

Correct Answer: D

Community vote distribution

**Comments**

ieffiong 3 months, 1 week ago

Selected Answer: D

When comparing the two:

Option C is less straightforward because State Manager isn't designed specifically for start/stop operations on RDS. It would likely require additional complexity (e.g., custom automation documents and potential workarounds) to achieve a scheduled start/stop functionality.

Option D leverages EventBridge's native scheduling capabilities to invoke Lambda functions, offering a more direct and streamlined approach with the least operational overhead.

upvoted 2 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: C

based on the AWS doc, you can use AWS Systems Manager State Manager associations to start and stop the RDS instances:
1. <https://docs.aws.amazon.com/systems-manager-automation-runbooks/latest/userguide/automation-ref-rds.html>

AWS-StartRdsInstance

AWS-StartStopAuroraCluster

AWS-StopRdsInstance

2. <https://docs.aws.amazon.com/systems-manager/latest/userguide/scheduling-automations-state-manager-associations.html>

D is also workable solution, but it will take more effort to implement.

upvoted 3 times

FlyingHawk 4 months, 3 weeks ago

If State Manager does not fully meet your needs, you can also use AWS Lambda with Amazon EventBridge to achieve similar functionality. This approach provides more flexibility but requires additional setup and management

<https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/automatically-stop-and-start-an-amazon-rds-db-instance-using-aws-systems-manager-maintenance-windows.html>

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: C

A - CloudWatch alarms are designed to respond to metric thresholds (like CPU/memory utilization) rather than schedules.

B - Trusted Advisor provides cost optimization and resource usage recommendations but can't schedule anything.

D - Works but more complicated than option C. We don't wanna manually compose and then maintain lambda functions right?

upvoted 1 times

Liongeek 5 months, 1 week ago

Selected Answer: C

Setting State Manager with AWS managed Documents (StartEC2Instance, StopEC2Instance, StartRDSInstance and StopRDSInstance) it's easier and faster than creating your own Lambda functions and event bridge rules.

Letter C is the answ.

upvoted 1 times

bujuman 6 months, 3 weeks ago

Selected Answer: D

AWS Systems Manager State Manager has limitations regarding MS SQL and the Db engine is not mentioned.

Check the this link to confirm D option : <https://aws.amazon.com/blogs/database/save-costs-by-automating-the-start-and-stop-of-amazon-rds-instances-with-aws-lambda-and-amazon-eventbridge/>

upvoted 4 times

Salilgen 4 months, 4 weeks ago

You are right: AWS Systems Manager has limitations (<https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/automatically-stop-and-start-an-amazon-rds-db-instance-using-aws-systems-manager-maintenance-windows.html#automatically-stop-and-start-an-amazon-rds-db-instance-using-aws-systems-manager-maintenance-windows-prereqs>).

Anyway, the question doesn't mention any of them

upvoted 1 times

ARV14 6 months, 1 week ago

It should be c as requirement is least operational overhead. Creating 4 lambda functions, event bridge rule, deploying through cloud formation works but adds operational overhead

upvoted 2 times

KaZimirovich 7 months, 2 weeks ago

Selected Answer: D

While AWS Systems Manager State Manager can be used to manage configuration states of AWS resources, it is generally more complex to set up for straightforward use cases like schedule-based starting and stopping of RDS instances compared to using a direct scheduling method through EventBridge.

upvoted 4 times

jingen11 7 months, 3 weeks ago

Selected Answer: C

Start, restart, or stop managed nodes and Amazon Relational Database Service (Amazon RDS) instances.

upvoted 1 times

hharbiordun85 8 months ago

D. Amazon EventBridge allows you to create rules based on a schedule (using cron expressions) to automate tasks. You can set up rules to start the RDS instances at the beginning of business hours and stop them at the end of business hours. By using AWS Lambda in conjunction with EventBridge, you can create functions that handle

upvoted 1 times

siheom 8 months, 1 week ago

Selected Answer: D

VOTE D

upvoted 2 times

AbhiBK 9 months, 1 week ago

To meet the requirement of running Amazon RDS DB instances only during business hours with the least operational overhead, the best solution would be:

D. Create an Amazon EventBridge rule that invokes AWS Lambda functions to start and stop the RDS instances. This approach allows you to automate the scheduling of start and stop actions using EventBridge rules, which can trigger Lambda functions based on a cron expression. This setup is straightforward and requires minimal ongoing management

upvoted 4 times

AbhiBK 9 months, 1 week ago

Option C, which involves using AWS Systems Manager State Manager to start and stop the RDS instances, is indeed a viable solution. It allows you to automate the process of keeping your RDS instances in a desired state, such as starting and stopping them on a schedule¹².

However, the reason Option D (using Amazon EventBridge with AWS Lambda) might be preferred for this scenario is due to its simplicity and flexibility. EventBridge rules can be easily configured with cron expressions to trigger Lambda functions, which can start and stop the RDS instances. This setup typically involves fewer steps and less configuration compared to setting up State Manager associations and IAM roles³.

Both options are valid, but Option D generally offers a more straightforward approach with potentially lower operational overhead

upvoted 2 times

dhewa 9 months, 3 weeks ago

Selected Answer: C

AWS Systems Manager State Manager allows you to automate the process of starting and stopping RDS instances based on a defined schedule.

upvoted 2 times

komorebi 10 months ago

Selected Answer: D

Answer is D

upvoted 4 times

nebajp 10 months ago

Selected Answer: C

Correct Answer us C - it allows you to define and automatically enforce desired configurations for EC2 and RDS.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #960

Topic 1

A consumer survey company has gathered data for several years from a specific geographic region. The company stores this data in an Amazon S3 bucket in an AWS Region.


The company has started to share this data with a marketing firm in a new geographic region. The company has granted the firm's AWS account access to the S3 bucket. The company wants to minimize the data transfer costs when the marketing firm requests data from the S3 bucket.

Which solution will meet these requirements?

- A. Configure the Requester Pays feature on the company's S3 bucket. **Most Voted**
- B. Configure S3 Cross-Region Replication (CRR) from the company's S3 bucket to one of the marketing firm's S3 buckets.
- C. Configure AWS Resource Access Manager to share the S3 bucket with the marketing firm AWS account.
- D. Configure the company's S3 bucket to use S3 Intelligent-Tiering Sync the S3 bucket to one of the marketing firm's S3 buckets.

Correct Answer: A

Community vote distribution



Comments

LeonSauveterre 5 months ago

Selected Answer: A

Read again - "The company wants to minimize the data transfer costs when the marketing firm requests data from the S3 bucket."

So as long as the company saves money, bravo. We don't care about the marketing firm. When enabling Requester Pays, the marketing firm (requester) pays for the costs of data requests and data transfer out of the S3 bucket, instead of the data owner (the consumer survey company).

B - CRR does not inherently reduce costs. In fact, replication incurs additional storage and transfer costs for the data owner, who would still bear the cost of replication, which contradicts the requirement to minimize their own data transfer costs.

C - RAM is used to share resources like S3 buckets, but it doesn't affect data transfer costs.

D - This helps reduce storage costs by moving objects between storage classes based on access patterns, but it doesn't reduce data transfer costs for requests made by the marketing firm.

data transfer costs for requests made by the marketing firm.

upvoted 1 times

Anyio 5 months, 2 weeks ago

Selected Answer: B

The correct answer is B. Configure S3 Cross-Region Replication (CRR) from the company's S3 bucket to one of the marketing firm's S3 buckets.

Explanation:

Option A: Incorrect. The Requester Pays feature can help distribute the cost of data requests by billing the requestor, but it does not reduce the overall data transfer costs. It only shifts who pays for the data transfer, not the amount.

Option B: Correct. S3 Cross-Region Replication (CRR) replicates the data to a bucket in the marketing firm's region. This allows the marketing firm to access the data from a local region, minimizing cross-region data transfer charges each time data is accessed.

upvoted 1 times

ARV14 6 months, 1 week ago

Selected Answer: B

This is the same question as question 88. One difference is "it's data transfer costs" in question 88. Here it is "minimize the data transfer costs". A or B can be correct based on whether they are asking for reducing data transfer costs for survey company or both of them.

upvoted 1 times

AMEJack 6 months, 1 week ago

Selected Answer: A

Sorry that was for the next question.

upvoted 1 times

martinadurcakova1 8 months ago

Selected Answer: A

A. Configure the Requester Pays feature on the company's S3 bucket.

Explanation:

A. Configuring the Requester Pays feature on the company's S3 bucket is the most appropriate solution. With Requester Pays, the marketing firm's AWS account will be responsible for the data transfer costs when accessing the data in the S3 bucket, minimizing the data transfer costs for the consumer survey company.

B. Configuring S3 Cross-Region Replication (CRR) from the company's S3 bucket to one of the marketing firm's S3 buckets would not be the most cost-effective solution, as the company would still be responsible for the data transfer costs.

upvoted 4 times

MatAlves 8 months, 2 weeks ago

The Requester Pays feature allows the bucket owner to offload the data transfer costs to the requester. When this feature is enabled, the marketing firm would pay for the data transfer when they access the data in the survey company's S3 bucket, which effectively minimizes costs for the survey company.

This means that the most cost-effective solution for the survey company, given that the marketing firm is accessing the data, is "A"

upvoted 2 times

[Removed] 9 months, 2 weeks ago

Answer B seems to be more logic, the question didn't mention which account will pay

upvoted 1 times

Abdullah2004 9 months, 1 week ago

question ask for method to reduce cost in survey company, so A will do it

upvoted 1 times

[Removed] 9 months, 3 weeks ago

Selected Answer: A

A sounds right
upvoted 2 times

JunsK1e 10 months ago

Selected Answer: A

Letter A
upvoted 2 times

pujithacg8 10 months ago

A is the correct answer
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/RequesterPaysBuckets.html>
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #961

Topic 1

A company uses AWS to host its public ecommerce website. The website uses an AWS Global Accelerator accelerator for traffic from the internet. The Global Accelerator accelerator forwards the traffic to an Application Load Balancer (ALB) that is the entry point for an Auto Scaling group.


The company recently identified a DDoS attack on the website. The company needs a solution to mitigate future attacks.

Which solution will meet these requirements with the LEAST implementation effort?

- A. Configure an AWS WAF web ACL for the Global Accelerator accelerator to block traffic by using rate-based rules
- B. Configure an AWS Lambda function to read the ALB metrics to block attacks by updating a VPC network ACL
- C. Configure an AWS WAF web ACL on the ALB to block traffic by using rate-based rules **Most Voted**
- D. Configure an Amazon CloudFront distribution in front of the Global Accelerator accelerator

Correct Answer: C

Community vote distribution



Comments

[Removed] **Highly Voted** 9 months, 3 weeks ago

Selected Answer: C

C as AWS WAF cannot be used directly on global accelerator
upvoted 9 times

dragosky 5 months, 3 weeks ago

it can be used, but to block HTTP/HTTPS requests.
upvoted 1 times

nebajp **Highly Voted** 10 months ago

Selected Answer: C

Correct answer C.
Global Accelerator does not work with WAF as it is suitable for TCP/UDP where as WAF is integrates with Application Load Balancer which is on Layer 7 on OSI model suitable for Web app (Http/Https)

balancer which is on Layer 7 on OSI model, suitable for web app (http/https)
upvoted 5 times

rpmaws Most Recent 8 months, 3 weeks ago

Selected Answer: C

WAF can be applied on ALB, API gateway or cloud front.
upvoted 5 times

sOI852POL 9 months ago

Selected Answer: C

Answer is C.
Note: AWS Global Accelerator itself doesn't support AWS WAF.

<https://repost.aws/knowledge-center/globalaccelerator-aws-waf-filter-layer7-traffic>
upvoted 5 times

dhewa 9 months, 3 weeks ago

Selected Answer: A

Configuring the WAF directly on the Global Accelerator ensures that malicious traffic is blocked before it reaches the Application Load Balancer, providing an additional layer of protection.
upvoted 2 times

officedepotadmin 9 months, 3 weeks ago

Selected Answer: A

Global Accelerator can be integrated with AWS WAF to provide protection at the edge, meaning malicious traffic can be blocked before it reaches your Application Load Balancer (ALB) or other resources in your AWS environment.
upvoted 2 times

AWS_Debu 9 months, 3 weeks ago

Answer is A

AWS Global Accelerator (GA) can be used with AWS Web Application Firewall (WAF) to protect applications from web exploits and DDoS attacks:

Block HTTP method and header attacks

GA, WAF, and the Application Load Balancer can block access to Layer 7 HTTP method and headers. WAF uses web access control list (web ACL) rules with the load balancer to evaluate incoming traffic and only forward requests that comply with the rules to the endpoint.

Detect and mitigate web application layer request floods

GA can protect web applications running on Application Load Balancer, and when used with WAF, it can also detect and mitigate web application layer request floods.

Prevent DDoS attacks

upvoted 1 times

GOTJ 4 months, 1 week ago

Unfortunately, AWS WAF cannot protect such a resource, according to this link:
<https://docs.aws.amazon.com/waf/latest/developerguide/waf-chapter.html>.

Currently supported resources are:

- o Amazon CloudFront distribution
- o Amazon API Gateway REST API
- o Application Load Balancer
- o AWS AppSync GraphQL API
- o Amazon Cognito user pool
- o AWS App Runner service
- o AWS Verified Access instance

upvoted 1 times

komorebi 10 months ago

Selected Answer: A

Answer is A
upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #962


Topic 1

A company uses an Amazon DynamoDB table to store data that the company receives from devices. The DynamoDB table supports a customer-facing website to display recent activity on customer devices. The company configured the table with provisioned throughput for writes and reads.

The company wants to calculate performance metrics for customer device data on a daily basis. The solution must have minimal effect on the table's provisioned read and write capacity.

Which solution will meet these requirements?

- A. Use an Amazon Athena SQL query with the Amazon Athena DynamoDB connector to calculate performance metrics on a recurring schedule.
- B. Use an AWS Glue job with the AWS Glue DynamoDB export connector to calculate performance metrics on a recurring schedule. **Most Voted**
- C. Use an Amazon Redshift COPY command to calculate performance metrics on a recurring schedule.
- D. Use an Amazon EMR job with an Apache Hive external table to calculate performance metrics on a recurring schedule.

Correct Answer: B*Community vote distribution***Comments****toyaji** **Highly Voted** 9 months ago**Selected Answer: B**

DynamoDB export connector literally "exports" table snapshot to s3 as dynamoDB-json object, then process on it. So it does not affect on read / write capacity on dynamoDB itself.

But Athena query directly on dynamoDB so affects on read / write capacity

upvoted 8 times

GOTJ **Most Recent** 4 months, 1 week ago**Selected Answer: B**

The trick here is to avoid, by any means, interact with the original DynamoDB table because doing so would affect the provisioned capacity. Option "B" is the only one that interacts with an EXPORTED COPY of the table. So, the question is: would

provisioned capacity. Option B is the only one that minimizes the impact on the table, so the question is: how provisioned capacity be affected by the DynamoDB export process? The answer is: NO!
(<https://docs.aws.amazon.com/glue/latest/dg/aws-glue-programming-etl-connect-dynamodb-home.html>):

"DynamoDB export to S3 allows you to export both full and incremental data from your DynamoDB table. Exports do not consume any read capacity units (RCUs) and have no impact on table performance and availability"

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: B

A - Athena allows you to use standard SQL, but it doesn't guarantee minimal impact on the table.
B - It minimizes direct read/write operations on the table because the exported data can be processed separately (for example, export to S3 then do the processing).
C - Cool but there's no native connector to directly pull data from DynamoDB. You must export the data to S3 or another source first, adding complexity.
D - EMR impacts throughput significantly.

upvoted 3 times

Anyio 5 months, 1 week ago

Selected Answer: B

The correct answer is B. Use an AWS Glue job with the AWS Glue DynamoDB export connector to calculate performance metrics on a recurring schedule.

Explanation:

Option B: Correct. AWS Glue offers an efficient method to extract, transform, and load (ETL) data from DynamoDB to Amazon S3, without affecting the table's provisioned throughput as significantly. Once data is in S3, it can be further processed or queried using AWS Glue or other analytics services, without impacting DynamoDB's immediate R/W provisioning.

upvoted 2 times

rockyykrish 6 months ago

Selected Answer: A

Explanation:

Amazon Athena with DynamoDB Connector:

Athena is a serverless interactive query service that allows you to run SQL queries directly on data in various storage systems, including DynamoDB, through the Athena DynamoDB connector.

This setup minimizes the impact on the provisioned read and write capacity of the DynamoDB table because the connector reads data efficiently without directly querying the table.

Performance Metrics Calculation:

Athena provides SQL capabilities to compute performance metrics on the data fetched via the DynamoDB connector. The queries can be scheduled using Amazon EventBridge or other automation tools to run daily.

Minimal Operational Overhead:

Athena requires no infrastructure to manage, making it a low-maintenance solution.

It is cost-effective since you pay only for the queries you run.

upvoted 1 times

tonybuivannghia 7 months ago

Selected Answer: A

I think A is correct. We can query the data from DynamoDB by Amazon Athena DynamoDB connector directly, not via S3 Bucket.

upvoted 2 times

Salilgen 4 months, 4 weeks ago

You can query the data from DynamoDB by Amazon Athena DynamoDB connector directly but this effect on the table's provisioned read and write capacity.

upvoted 1 times

tonybuivannghia 7 months ago

<https://docs.aws.amazon.com/athena/latest/ug/connectors-dynamodb.html>

upvoted 1 times

snover 8 months, 1 week ago

spvrcu 8 months, 1 week ago

Selected Answer: B

<https://aws.amazon.com/blogs/big-data/accelerate-amazon-dynamodb-data-access-in-aws-glue-jobs-using-the-new-aws-glue-dynamodb-elt-connector/>

upvoted 3 times

JoeTromundo 8 months, 2 weeks ago

Selected Answer: B

The DynamoDB export connector allows you to export data from DynamoDB to other storage solutions like Amazon S3 without consuming the table's provisioned read capacity, ensuring minimal impact on the performance of the table.

upvoted 3 times

Bogey 8 months, 4 weeks ago

B.
Instead, the new AWS Glue DynamoDB export connector reads DynamoDB data from the snapshot, which is exported from DynamoDB tables. This approach has following benefits:

It doesn't consume read capacity units of the source DynamoDB tables

upvoted 2 times

AbhiBK 9 months, 1 week ago

To calculate performance metrics for customer device data on a daily basis with minimal effect on the table's provisioned read and write capacity, the best solution would be:

A. Use an Amazon Athena SQL query with the Amazon Athena DynamoDB connector to calculate performance metrics on a recurring schedule. This approach allows you to run SQL queries directly on the data stored in DynamoDB without impacting the provisioned throughput, as Athena queries are serverless and do not consume DynamoDB read or write capacity1.

upvoted 2 times

RealPro111 9 months, 2 weeks ago

Selected Answer: B

The right answer is B

upvoted 3 times

siheom 9 months, 2 weeks ago

VOTE B

upvoted 2 times

ksdpmx 9 months, 2 weeks ago

why is B wrong.. Glue DynamoDB export connector will read data from PITR instead of DynamoDB directly..

upvoted 2 times

dhewa 9 months, 3 weeks ago

I go with A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #963

Topic 1

A solutions architect is designing the cloud architecture for a new stateless application that will be deployed on AWS. The solutions architect created an Amazon Machine Image (AMI) and launch template for the application.

Based on the number of jobs that need to be processed, the processing must run in parallel while adding and removing application Amazon EC2 instances as needed. The application must be loosely coupled. The job items must be durably stored.

Which solution will meet these requirements?

- A. Create an Amazon Simple Notification Service (Amazon SNS) topic to send the jobs that need to be processed. Create an Auto Scaling group by using the launch template with the scaling policy set to add and remove EC2 instances based on CPU usage.
- B. Create an Amazon Simple Queue Service (Amazon SQS) queue to hold the jobs that need to be processed. Create an Auto Scaling group by using the launch template with the scaling policy set to add and remove EC2 instances based on network usage.
- C. Create an Amazon Simple Queue Service (Amazon SQS) queue to hold the jobs that need to be processed. Create an Auto Scaling group by using the launch template with the scaling policy set to add and remove EC2 instances based on the number of items in the SQS queue. **Most Voted**
- D. Create an Amazon Simple Notification Service (Amazon SNS) topic to send the jobs that need to be processed. Create an Auto Scaling group by using the launch template with the scaling policy set to add and remove EC2 instances based on the number of messages published to the SNS topic.

Correct Answer: C*Community vote distribution*C (100%)**Comments****LeonSauveterre** 5 months ago**Selected Answer: C**

SQS for sure, but I actually swayed between B and C. After a while, I can say now I'm sure that scaling based on network usage doesn't directly have anything to do with the number of jobs in the queue. Or worse, this could result in burst-caused downtime or other poor scaling decisions.

upvoted 3 times

dhewa 9 months, 3 weeks ago

Answer is C: SQS is your first cue then scaling based on the number of requests

upvoted 3 times

[Removed] 9 months, 3 weeks ago

Selected Answer: C

"Based on the number of jobs that need to be processed"

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #964

Topic 1

A global ecommerce company uses a monolithic architecture. The company needs a solution to manage the increasing volume of product data. The solution must be scalable and have a modular service architecture. The company needs to maintain its structured database schemas. The company also needs a storage solution to store product data and product images.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use an Amazon EC2 instance in an Auto Scaling group to deploy a containerized application. Use an Application Load Balancer to distribute web traffic. Use an Amazon RDS DB instance to store product data and product images.
- B. Use AWS Lambda functions to manage the existing monolithic application. Use Amazon DynamoDB to store product data and product images. Use Amazon Simple Notification Service (Amazon SNS) for event-driven communication between the Lambda functions.
- C. Use Amazon Elastic Kubernetes Service (Amazon EKS) with an Amazon EC2 deployment to deploy a containerized application. Use an Amazon Aurora cluster to store the product data. Use AWS Step Functions to manage workflows. Store the product images in Amazon S3 Glacier Deep Archive.
- D. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate to deploy a containerized application. Use Amazon RDS with a Multi-AZ deployment to store the product data. Store the product images in an Amazon S3 bucket.

Most Voted**Correct Answer: D***Community vote distribution*

D (93%)

C (7%)

Comments**LeonSauveterre** **Highly Voted** 5 months ago**Selected Answer: D**

Based on my personal experience, "monolith" or "monolithic" is almost always ECS + Fargate. For your reference only ;) upvoted 6 times

trinh_le **Most Recent** 5 months, 1 week ago**Selected Answer: C**

Keyword scalable. So must use serverless like lambda or fargate (B or D)

Keyword modular => ECS => D
upvoted 1 times

GOTJ 4 months, 1 week ago

I think you are taking "modular" in the wrong way. I think they mean "modular" as an alternative to "monolithic". In other words, "modular" --> "Micro services" --> ECS/EKS
upvoted 1 times

martinadurcakova1 8 months ago

Selected Answer: D

D is right
upvoted 2 times

dhewa 9 months, 3 weeks ago

Selected Answer: D

D gives us a server-less solution which is what we want.
upvoted 4 times

[Removed] 9 months, 3 weeks ago

Selected Answer: D

D sounds right
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #965

Topic 1

A company is migrating an application from an on-premises environment to AWS. The application will store sensitive data in Amazon S3. The company must encrypt the data before storing the data in Amazon S3.

Which solution will meet these requirements?

- A. Encrypt the data by using client-side encryption with customer managed keys. **Most Voted**
- B. Encrypt the data by using server-side encryption with AWS KMS keys (SSE-KMS).
- C. Encrypt the data by using server-side encryption with customer-provided keys (SSE-C).
- D. Encrypt the data by using client-side encryption with Amazon S3 managed keys.

Correct Answer: A*Community vote distribution***Comments****Cpso** 6 months, 1 week ago**Selected Answer: A**

Question is unclear. "before" store= before send api or before save?

upvoted 2 times

Tygrins 7 months, 2 weeks ago

To meet the requirement of encrypting sensitive data before storing it in Amazon S3, the company should use Server-Side Encryption with Customer-Provided Keys (SSE-C) or Client-Side Encryption. However, the most common and effective approach in AWS is Server-Side Encryption with AWS Key Management Service (SSE-KMS).

upvoted 1 times

JoeTromundo 8 months, 2 weeks ago**Selected Answer: A**

If the answer is B, how will the data be encrypted before being stored in Amazon S3? It has to be client-side encryption.

upvoted 2 times

dhewa 9 months, 3 weeks ago**Selected Answer: B**

AWS handles the encryption and decryption process, simplifying the management of encryption keys.

upvoted 1 times

[Removed] 9 months, 3 weeks ago

Selected Answer: A

If client-side encryption is used, the keys must be managed by the customer.

upvoted 3 times

George1990 10 months ago

Correct is A

upvoted 3 times

JunsK1e 10 months, 1 week ago

Selected Answer: A

Before you store the data in S3 you can use client side encryption

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #966

Topic 1

A company wants to create an Amazon EMR cluster that multiple teams will use. The company wants to ensure that each team's big data workloads can access only the AWS services that each team needs to interact with. The company does not want the workloads to have access to Instance Metadata Service Version 2 (IMDSv2) on the cluster's underlying EC2 instances.

Which solution will meet these requirements?

- A. Configure interface VPC endpoints for each AWS service that the teams need. Use the required interface VPC endpoints to submit the big data workloads.
- B. Create EMR runtime roles. Configure the cluster to use the runtime roles. Use the runtime roles to submit the big data workloads. **Most Voted**
- C. Create an EC2 IAM instance profile that has the required permissions for each team. Use the instance profile to submit the big data workloads.
- D. Create an EMR security configuration that has the EnableApplicationScopedIAMRole option set to false. Use the security configuration to submit the big data workloads.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**FlyingHawk** 4 months, 3 weeks ago**Selected Answer: B**

When you specify a runtime role for an Amazon EMR step, the jobs or queries that you submit can only access AWS resources that the policies attached to the runtime role allow. These jobs and queries can't access the Instance Metadata Service on the EC2 instances of the cluster or use the EC2 instance profile of the cluster to access any AWS resources.

<https://aws.amazon.com/blogs/big-data/introducing-runtime-roles-for-amazon-emr-steps-use-iam-roles-and-aws-lake-formation-for-access-control-with-amazon-emr/>

<https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-steps-runtime-roles.html>

upvoted 3 times

EllenLiu 5 months, 2 weeks ago**Selected Answer: B**

runtime role is used to resolve the issue that there is only union permissions for all jobs and queries we can use. just like it mentions, for ec2, IMDSv2 is mandatory, however we can use runtime role to access instance metadata on ec2 avoid using IMDSv2

upvoted 1 times

martinadurcakova1 8 months ago

Selected Answer: B

B. Creating EMR runtime roles and configuring the cluster to use them is the correct solution. EMR runtime roles allow you to grant specific permissions to the big data workloads, ensuring that each team's workloads can only access the required AWS services. Additionally, the runtime roles can be configured to disable access to IMDSv2, meeting the requirement.

upvoted 2 times

dhewa 9 months, 3 weeks ago

Selected Answer: B

This approach avoids the need for workloads to access the Instance Metadata Service (IMDSv2) on the underlying EC2 instances, as the permissions are managed through the runtime roles.

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: B

Explanation:

EMR Runtime Roles: By creating EMR runtime roles, you can assign specific IAM roles to individual EMR jobs or steps. Each role can have fine-grained permissions, allowing you to restrict access to only the AWS services each team needs. This provides a highly controlled environment where each team's workload operates under the principle of least privilege.

IMDSv2 Access: When using runtime roles, you do not rely on the EC2 instance profile for service access, thereby minimizing the need for the workloads to access the Instance Metadata Service. This can help in reducing the risk of unauthorized access to IMDSv2.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #967

Topic 1

A solutions architect is designing an application that helps users fill out and submit registration forms. The solutions architect plans to use a two-tier architecture that includes a web application server tier and a worker tier.

The application needs to process submitted forms quickly. The application needs to process each form exactly once. The solution must ensure that no data is lost.

Which solution will meet these requirements?

- A. Use an Amazon Simple Queue Service (Amazon SQS) FIFO queue between the web application server tier and the worker tier to store and forward form data. **Most Voted**
- B. Use an Amazon API Gateway HTTP API between the web application server tier and the worker tier to store and forward form data.
- C. Use an Amazon Simple Queue Service (Amazon SQS) standard queue between the web application server tier and the worker tier to store and forward form data.
- D. Use an AWS Step Functions workflow. Create a synchronous workflow between the web application server tier and the worker tier that stores and forwards form data.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**blehbleh** 8 months, 1 week ago**Selected Answer: A**

Unlike standard queues, FIFO queues don't introduce duplicate messages. FIFO queues help you avoid sending duplicates to a queue. If you retry the SendMessage action within the 5-minute deduplication interval, Amazon SQS doesn't introduce any duplicates into the queue.

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/FIFO-queues-exactly-once-processing.html>
upvoted 2 times

dhewa 9 months, 3 weeks ago**Selected Answer: A**

Selected Answer: A

Keywords in the qsn (quickly and once) = SQS FIFO
upvoted 3 times

[Removed] 9 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 2 times

komorebi 10 months, 1 week ago

Selected Answer: A

Answer is A
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #968

Topic 1

A finance company uses an on-premises search application to collect streaming data from various producers. The application provides real-time updates to search and visualization features.

The company is planning to migrate to AWS and wants to use an AWS native solution.

Which solution will meet these requirements?

- A. Use Amazon EC2 instances to ingest and process the data streams to Amazon S3 buckets for storage. Use Amazon Athena to search the data. Use Amazon Managed Grafana to create visualizations.
- B. Use Amazon EMR to ingest and process the data streams to Amazon Redshift for storage. Use Amazon Redshift Spectrum to search the data. Use Amazon QuickSight to create visualizations.
- C. Use Amazon Elastic Kubernetes Service (Amazon EKS) to ingest and process the data streams to Amazon DynamoDB for storage. Use Amazon CloudWatch to create graphical dashboards to search and visualize the data.
- D. Use Amazon Kinesis Data Streams to ingest and process the data streams to Amazon OpenSearch Service. Use OpenSearch Service to search the data. Use Amazon QuickSight to create visualizations. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

Gbemi **Highly Voted** 10 months, 1 week ago

Answer is D
using Amazon Kinesis Data Streams, Amazon OpenSearch Service, and Amazon QuickSight provides a comprehensive and AWS-native solution that meets the requirements of real-time data ingestion, search, and visualization
upvoted 6 times

dhewa **Most Recent** 9 months, 3 weeks ago

Selected Answer: D

Keyword "streaming" hence Amazon Kinesis Data Streams as it is designed for real-time data ingestion and processing, making it ideal for collecting streaming data from various producers.
upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: D

D sounds right
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #969

Topic 1

A company currently runs an on-premises application that uses ASP.NET on Linux machines. The application is resource-intensive and serves customers directly.

The company wants to modernize the application to .NET. The company wants to run the application on containers and to scale based on Amazon CloudWatch metrics. The company also wants to reduce the time spent on operational maintenance activities.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use AWS App2Container to containerize the application. Use an AWS CloudFormation template to deploy the application to Amazon Elastic Container Service (Amazon ECS) on AWS Fargate. **Most Voted**
- B. Use AWS App2Container to containerize the application. Use an AWS CloudFormation template to deploy the application to Amazon Elastic Container Service (Amazon ECS) on Amazon EC2 instances.
- C. Use AWS App Runner to containerize the application. Use App Runner to deploy the application to Amazon Elastic Container Service (Amazon ECS) on AWS Fargate.
- D. Use AWS App Runner to containerize the application. Use App Runner to deploy the application to Amazon Elastic Kubernetes Service (Amazon EKS) on Amazon EC2 instances.

Correct Answer: A

Community vote distribution

A (100%)

Comments

LeonSauveterre 5 months ago

Selected Answer: A

App2Container: A tool to automatically containerize .NET applications.

Fargate: A serverless compute engine for containers, eliminating the need to manage underlying EC2 instances. It integrates with CloudWatch for scaling based on custom metrics.

App Runner: A fully managed service designed for simple application deployment and scaling, but it also deploys applications *directly*, which means it doesn't deploy to ECS or EKS.

upvoted 3 times

JoeTromundo 8 months, 2 weeks ago

Selected Answer: A

Meets all requirements.
upvoted 2 times

dhewa 9 months, 3 weeks ago

Selected Answer: A

A gives us a serverless solution
upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: A

AWS App2Container: This service helps you easily containerize existing applications, such as your ASP.NET application, reducing the complexity of the containerization process.

Amazon ECS on AWS Fargate: Fargate is a serverless compute engine for containers that eliminates the need to manage the underlying EC2 instances, significantly reducing operational overhead. You only need to focus on your containerized application, while AWS handles the infrastructure.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #970

Topic 1

A company is designing a new internal web application in the AWS Cloud. The new application must securely retrieve and store multiple employee usernames and passwords from an AWS managed service.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Store the employee credentials in AWS Systems Manager Parameter Store. Use AWS CloudFormation and the BatchGetSecretValue API to retrieve usernames and passwords from Parameter Store.
- B. Store the employee credentials in AWS Secrets Manager. Use AWS CloudFormation and AWS Batch with the BatchGetSecretValue API to retrieve the usernames and passwords from Secrets Manager.
- C. Store the employee credentials in AWS Systems Manager Parameter Store. Use AWS CloudFormation and AWS Batch with the BatchGetSecretValue API to retrieve the usernames and passwords from Parameter Store.
- D. Store the employee credentials in AWS Secrets Manager. Use AWS CloudFormation and the BatchGetSecretValue API to retrieve the usernames and passwords from Secrets Manager. **Most Voted**

Correct Answer: D*Community vote distribution*

D (100%)

Comments**JoeTromundo** **Highly Voted** 8 months, 2 weeks ago**Selected Answer: D**

Option D is similar to option B, but option B unnecessarily introduces AWS Batch into the solution. AWS Batch is designed for executing batch jobs and is not required for the use case of retrieving secrets in a web application. This adds complexity and overhead without benefit.

upvoted 8 times

FlyingHawk **Most Recent** 4 months, 2 weeks ago**Selected Answer: D**

BatchGetSecretValue API is with Secrets Manager, not related to AWS Batch, it allows you to retrieve 1 to 20 secrets via a single API call.

<https://aws.amazon.com/blogs/security/how-to-use-the-batchgetsecretsvalue-api-to-improve-your-client-side-applications-with-aws-secrets-manager/>

upvoted 2 times

upvoted 2 times

spoved 8 months, 1 week ago

Selected Answer: D

https://docs.aws.amazon.com/secretsmanager/latest/apireference/API_BatchGetSecretValue.html

upvoted 4 times

viejito 8 months, 4 weeks ago

respuesta correcta D : La opción D proporciona una solución adecuada y segura para gestionar y recuperar secretos, con un enfoque en la menor sobrecarga operativa posible al utilizar AWS Secrets Manager junto con la API correcta y la integración con AWS CloudFormation.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #971


Topic 1

A company that is in the ap-northeast-1 Region has a fleet of thousands of AWS Outposts servers. The company has deployed the servers at remote locations around the world. All the servers regularly download new software versions that consist of 100 files. There is significant latency before all servers run the new software versions.

The company must reduce the deployment latency for new software versions.

Which solution will meet this requirement with the LEAST operational overhead?

- A. Create an Amazon S3 bucket in ap-northeast-1. Set up an Amazon CloudFront distribution in ap-northeast-1 that includes a CachingDisabled cache policy. Configure the S3 bucket as the origin. Download the software by using signed URLs.
- B. Create an Amazon S3 bucket in ap-northeast-1. Create a second S3 bucket in the us-east-1 Region. Configure replication between the buckets. Set up an Amazon CloudFront distribution that uses ap-northeast-1 as the primary origin and us-east-1 as the secondary origin. Download the software by using signed URLs.
- C. Create an Amazon S3 bucket in ap-northeast-1. Configure Amazon S3 Transfer Acceleration. Download the software by using the S3 Transfer Acceleration endpoint.
- D. Create an Amazon S3 bucket in ap-northeast-1. Set up an Amazon CloudFront distribution. Configure the S3 bucket as the origin. Download the software by using signed URLs. **Most Voted**

Correct Answer: D*Community vote distribution***Comments****JoeTromundo** **Highly Voted** 8 months, 2 weeks ago**Selected Answer: D**

By setting up a CloudFront distribution with the S3 bucket as the origin, the software versions can be cached at edge locations close to the remote AWS Outposts servers. This reduces the latency of downloading new software versions because the servers can access the content from the nearest CloudFront edge location instead of downloading it directly from the S3 bucket in the ap-northeast-1 Region.

About option C: S3 Transfer Acceleration optimizes the upload and download of files to S3 by routing through optimized network paths. However, it is primarily designed to improve performance when transferring data over long distances to S3 (uploads). It does not provide the same level of global caching and latency reduction as CloudFront for large-scale distribution.

upvoted 7 times

ashishs174 Highly Voted 8 months, 2 weeks ago

Option C is correct.

Option D brings additional overhead in terms of setting up, so opted out.

Also from google : If you want to download software quickly from a large distance, you should generally use S3 Transfer Acceleration to download from your S3 bucket, as it leverages Amazon's global edge network to significantly speed up long-distance data transfers, making it ideal for large file downloads across continents; while CloudFront is better suited for delivering smaller content like website assets to users globally, where faster response times and caching are critical

upvoted 6 times

FlyingHawk Most Recent 4 months, 1 week ago

Selected Answer: D

Using Amazon CloudFront as a global content delivery network (CDN) caches software updates closer to the remote Outposts servers at AWS edge locations. This significantly reduces latency compared to downloading files directly from S3.

- Global Acceleration – CloudFront caches the software in edge locations worldwide, reducing latency for remote locations.
- Efficient Distribution – Servers in different regions fetch files from the closest CloudFront edge cache instead of always pulling from ap-northeast-1.
- Signed URLs for Security – Ensures only authorized servers can download the software.
- Low Operational Overhead – No need for replication, manual distribution, or complex networking setups.

upvoted 1 times

FlyingHawk 4 months, 1 week ago

C. S3 Transfer Acceleration □ Improves upload speed but does not help with global caching. Each Outpost server still downloads directly from ap-northeast-1, causing high latency.

upvoted 1 times

dfgdsfgfdgreg 4 months, 2 weeks ago

Selected Answer: C

Including "global" makes CloudFront stand out as the obvious choice. But S3 global acceleration requires the least operational overhead while also assisting with latency issues. The mention of hundreds of files also points towards s3 acceleration.

upvoted 1 times

Saliigen 4 months, 4 weeks ago

Selected Answer: C

IMO answer is C.

A is out because disabling caching, it involves higher latency than C

B is out because it doesn't meet the requirements: adds complexity without significant latency reduction

D meets the requirements but it requires more operational overhead than C

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: D

A - CachingDisabled is the polar opposite of our goal.

B - CloudFront with dual-region origins adds unnecessary complexity without significant latency reduction.

C - Faster upload/download speeds compared to direct S3 access, but the distribution latency is still there!

D - Looks good. Just one small thing: files may take a little longer to propagate to edge locations initially, but this can be mitigated by configuring an appropriate cache invalidation policy. Once cached at edge locations, subsequent downloads would be much faster.

upvoted 1 times

deacon967 5 months, 2 weeks ago

Selected Answer: C

C : all other options use CloudFront. CloudFront is a CDN, not needed for software download.

upvoted 1 times

RamanAgarwal 5 months, 3 weeks ago

Selected Answer: C

Should use S3 Transfer acceleration to avoid CF setup overhead. It will help with improving download speed

upvoted 1 times

JA2018 5 months, 3 weeks ago

Selected Answer: A

What about A?

1. Minimized operational overhead:

This option provides the least operational overhead because it only requires setting up a single S3 bucket and a CloudFront distribution in the same region (ap-northeast-1), with a simple configuration to disable caching on the CloudFront distribution, ensuring fresh software updates are always delivered to Outposts servers.

2. Reduced latency:

By using a CloudFront distribution with a CachingDisabled policy, the software updates will be delivered from the closest edge location to each Outpost server, significantly reducing latency due to geographical proximity.

3. Signed URLs for security:

Using signed URLs ensures that only authorized users can access the software updates, enhancing security.

upvoted 1 times

dhewa 9 months, 3 weeks ago

Selected Answer: D

This approach leverages the global presence of CloudFront to minimize latency and simplifies the deployment process, reducing operational overhead.

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: D

S3 Transfer Acceleration speeds up uploads and downloads to S3 over long distances but does not provide the global caching benefits that CloudFront offers. CloudFront is more suitable for reducing latency when distributing content globally.

upvoted 4 times

nebajp 10 months ago

Selected Answer: C

Correct answer is C

Because for new software version we do not use Cloudfront,

upvoted 2 times

pelrock 10 months ago

Selected Answer: D

The question asks for the solution with the least operational overhead. Option D has minimal setup and maintenance requirements, as CloudFront is a managed service that handles caching and distribution across the globe, reducing latency without requiring additional configurations or services.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #972

Topic 1

A company currently runs an on-premises stock trading application by using Microsoft Windows Server. The company wants to migrate the application to the AWS Cloud.

The company needs to design a highly available solution that provides low-latency access to block storage across multiple Availability Zones.


Which solution will meet these requirements with the LEAST implementation effort?

A. Configure a Windows Server cluster that spans two Availability Zones on Amazon EC2 instances. Install the application on both cluster nodes. Use Amazon FSx for Windows File Server as shared storage between the two cluster nodes.

B. Configure a Windows Server cluster that spans two Availability Zones on Amazon EC2 instances. Install the application on both cluster nodes. Use Amazon Elastic Block Store (Amazon EBS) General Purpose SSD (gp3) volumes as storage attached to the EC2 instances. Set up application-level replication to sync data from one EBS volume in one Availability Zone to another EBS volume in the second Availability Zone.

C. Deploy the application on Amazon EC2 instances in two Availability Zones. Configure one EC2 instance as active and the second EC2 instance in standby mode. Use an Amazon FSx for NetApp ONTAP Multi-AZ file system to access the data by using Internet Small Computer Systems Interface (iSCSI) protocol. **Most Voted**

D. Deploy the application on Amazon EC2 instances in two Availability Zones. Configure one EC2 instance as active and the second EC2 instance in standby mode. Use Amazon Elastic Block Store (Amazon EBS) Provisioned IOPS SSD (io2) volumes as storage attached to the EC2 instances. Set up Amazon EBS level replication to sync data from one io2 volume in one Availability Zone to another io2 volume in the second Availability Zone.

Correct Answer: C*Community vote distribution***Comments****LeonSauveterre** **Highly Voted** 5 months ago**Selected Answer: C**

Keep in mind: WE NEED LOW LATENCY BLOCK STORAGE.

A - FSx for Windows File Server is not designed for block storage. It is optimized for shared file storage.
B - God this requires high implementation effort.
C - FSx for NetApp ONTAP supports block storage (via iSCSI) and provides Multi-AZ fault tolerance out of the box.
D - EBS io2 volumes with replication across AZs are hard and complex to configure and manage. Even if you think this is acceptable, still it's definitely much harder than option C.

upvoted 5 times

LeonSauveterre 5 months ago

I know it is almost correct to look for "FSx for Windows" when you spot "Microsoft Windows Server" in the question, but that's far from enough to get the correct answer.

upvoted 1 times

Anyio Most Recent 5 months, 2 weeks ago

Selected Answer: C

The correct answer is C.

Explanation:

Option C: Correct. Amazon FSx for NetApp ONTAP provides a managed file storage solution that supports iSCSI protocol for block storage access, enabling low-latency, highly available access to shared storage across multiple Availability Zones. This solution requires less implementation effort because Amazon FSx manages the complex underlying multi-AZ configuration and data replication automatically.

Option A: Incorrect. While Amazon FSx for Windows File Server provides a managed shared storage solution, it is primarily used for SMB-based file shares rather than block storage needs like those typically demanded by stock trading applications. Also, it might introduce additional latency for block-level access.

upvoted 3 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: C

block storage => FSx for NetApp ONTAP not only provides iSCSI(for block) support, but also offers comprehensive protocol support including SMB and NFS.

upvoted 2 times

ARV14 6 months, 1 week ago

Selected Answer: C

Fsx windows file server does not provide block storage, also Ontap provides very low latency and provides various compatible storage including block

upvoted 1 times

rosanna 6 months, 1 week ago

Selected Answer: D

block means EBS, and I've searched but I don't see Fsx supporting shared block storage, not that I know of, and thinking about if it did exist, I would cost more than a replica- both operationally and money wise. And mostly we don't need complex methods so I'd go for D

upvoted 1 times

ckhemani 6 months, 1 week ago

Selected Answer: A

Answer A. Although they are looking for Block storage which is EBS but EBS can't be shared across two availability zones

.
Multi-Attach feature: While you can use the "Multi-Attach" feature to attach a single EBS volume to multiple instances within the same Availability Zone, this does not allow sharing across different Availability Zones.

Configuring NETAPP is not least operational.

upvoted 1 times

Bwhizzy 7 months, 3 weeks ago

Selected Answer: A

Answer is A.

Amazon FSx for Windows File Server provides fully managed shared storage designed for Windows workloads. It offers seamless integration with Windows applications and supports Windows Server Failover Clustering (WSFC), making it the ideal

seamless integration with windows applications and supports windows server failover clustering (WSFC), making it the ideal choice for clustering applications.
High availability is provided as the cluster spans multiple Availability Zones, ensuring that the application continues to function in case of an AZ failure.
Low-latency access to shared block storage is achievable with FSx, and it reduces the complexity compared to setting up replication between multiple Amazon EBS volumes.
Least implementation effort: You don't need to manage block-level replication or create complex replication setups between AZs, as FSx for Windows File Server handles this automatically. It simplifies storage management and offers native Windows support.

upvoted 2 times

martinadurcakova1 8 months ago

Using Amazon FSx for Windows File Server would not be the best solution, as the company requires block storage, not file storage.

D. Deploying the application on EC2 instances and using Amazon EBS Provisioned IOPS SSD (io2) volumes with EBS-level replication is the solution with the least implementation effort for the following reasons:

EBS volumes provide the required block storage for the application.

Using EBS-level replication to sync data between Availability Zones is a built-in feature, requiring less implementation effort compared to setting up custom replication mechanisms.

Provisioned IOPS SSD (io2) volumes ensure the low-latency access to the block storage required by the application.

By using Amazon EBS Provisioned IOPS SSD (io2) volumes with EBS-level replication, the company can meet the requirements for high availability and low-latency access to block storage with the least implementation effort.

upvoted 4 times

spoved 8 months, 1 week ago

Selected Answer: D

block storage -> B or D

io1, io2 support EBS Multi-Attach -> D

upvoted 3 times

HappyG 8 months, 1 week ago

Selected Answer: C

Key word in the question is block storage which eliminates A as a possibility. FSx for NetApp ONTAP provides low-latency block storage, multi-AZ high availability, and requires the least implementation effort because it is a fully managed service with built-in replication and failover capabilities. This makes it the most suitable and cost-effective option for the company's needs.

upvoted 2 times

trongod05 8 months ago

Amazon FSx for NetApp ONTAP also provides shared block storage over the iSCSI and NVMe-over-TCP protocols. Seems to be a complete match there.

upvoted 2 times

JoeTromundo 8 months, 2 weeks ago

Selected Answer: A

For those who are choosing option other than A: what part of "with the LEAST implementation effort" do you not understand?

upvoted 2 times

HappyG 8 months, 1 week ago

The part where it calls for block storage.

upvoted 2 times

ashishs174 8 months, 2 weeks ago

Answer is B

EFS is not Block Storage, EBS is.(Thus, A and C are out)

upvoted 1 times

rpmaws 8 months, 3 weeks ago

Selected Answer: C

due to this line in question "low-latency access to block storage" C looks appropriate.

upvoted 3 times

toyaji 9 months ago

Selected Answer: A

A - Windows cluster across two availability zones satisfies highly available condition and FSx satisfies storage access from multiple ec2 especially for Windows server.

B - "Set up application-level replication to sync data" not for LEAST effort condition

C, D - "in standby mode" need to be activated manually so its not highly available

upvoted 2 times

elmyth 9 months, 1 week ago

Selected Answer: B

Block storage

upvoted 2 times

elmyth 9 months, 1 week ago

Changed for C, bc application-level replication is more difficult to implement.

upvoted 2 times

dhewa 9 months, 3 weeks ago

Selected Answer: A

Windows = FSX

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #973

Topic 1

A company is designing a web application with an internet-facing Application Load Balancer (ALB).


The company needs the ALB to receive HTTPS web traffic from the public internet. The ALB must send only HTTPS traffic to the web application servers hosted on the Amazon EC2 instances on port 443. The ALB must perform a health check of the web application servers over HTTPS on port 8443.

Which combination of configurations of the security group that is associated with the ALB will meet these requirements? (Choose three.)

- A. Allow HTTPS inbound traffic from 0.0.0.0/0 for port 443. **Most Voted**
- B. Allow all outbound traffic to 0.0.0.0/0 for port 443.
- C. Allow HTTPS outbound traffic to the web application instances for port 443. **Most Voted**
- D. Allow HTTPS inbound traffic from the web application instances for port 443.
- E. Allow HTTPS outbound traffic to the web application instances for the health check on port 8443. **Most Voted**
- F. Allow HTTPS inbound traffic from the web application instances for the health check on port 8443.

Correct Answer: ACE

Community vote distribution

**Comments**

spoved 8 months, 1 week ago

Selected Answer: ACE

<https://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-update-security-groups.html>

The following rules are recommended for an internet-facing load balancer.

upvoted 4 times

JoeTromundo 8 months, 2 weeks ago

Selected Answer: ACE

This question is poorly worded: assuming that, by default in security groups, all OUTBOUND connections are ALLOWED and all INBOUND connections are DENIED, options C and E would not even need to be configured. What would be needed is to create

a security group for the EC2 instances allowing INBOUND connections from the ALB security group to the EC2 instances security group on ports 443 and 8443.

upvoted 3 times

dhewa 9 months, 3 weeks ago

Selected Answer: ACE

- A. This allows the ALB to receive HTTPS traffic from the public internet.
- C. This ensures that the ALB can send HTTPS traffic to the web application servers.
- E. This allows the ALB to perform health checks on the web application servers over HTTPS on port 8443.

upvoted 4 times

dhewa 9 months, 3 weeks ago

Selected Answer: ADE

- A. This allows the ALB to receive HTTPS traffic from the public internet.
- C. This ensures that the ALB can send HTTPS traffic to the web application servers.
- E. This allows the ALB to perform health checks on the web application servers over HTTPS on port 8443.

upvoted 1 times

aragon_saa 9 months, 3 weeks ago

Selected Answer: ACE

Answer is ACE

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: ACE

A. Allow HTTPS inbound traffic from 0.0.0.0/0 for port 443.

This allows the ALB to receive HTTPS traffic from the public internet on port 443.

C. Allow HTTPS outbound traffic to the web application instances for port 443.

This allows the ALB to forward HTTPS traffic to the web application servers on port 443.

E. Allow HTTPS outbound traffic to the web application instances for the health check on port 8443.

This allows the ALB to perform health checks on the web application servers over HTTPS on port 8443.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #974

Topic 1

A company hosts an application on AWS. The application gives users the ability to upload photos and store the photos in an Amazon S3 bucket. The company wants to use Amazon CloudFront and a custom domain name to upload the photo files to the S3 bucket in the eu-west-1 Region.

Which solution will meet these requirements? (Choose two.)

A. Use AWS Certificate Manager (ACM) to create a public certificate in the us-east-1 Region. Use the certificate in CloudFront. **Most Voted**

B. Use AWS Certificate Manager (ACM) to create a public certificate in eu-west-1. Use the certificate in CloudFront.


C. Configure Amazon S3 to allow uploads from CloudFront. Configure S3 Transfer Acceleration.

D. Configure Amazon S3 to allow uploads from CloudFront origin access control (OAC). **Most Voted**

E. Configure Amazon S3 to allow uploads from CloudFront. Configure an Amazon S3 website endpoint.

Correct Answer: AD

Community vote distribution



Comments

JoeTromundo **Highly Voted** 8 months, 2 weeks ago

Selected Answer: AD

Amazon CloudFront requires an SSL/TLS certificate to use HTTPS with a custom domain name. This certificate **MUST** be provisioned in the us-east-1 Region, regardless of where your content is hosted. This is because CloudFront only supports certificates in the us-east-1 Region for use with custom domain names.

Origin Access Control (OAC) is a feature that allows you to securely upload content to an S3 bucket using CloudFront. It provides fine-grained access control and ensures that only CloudFront can upload files to the S3 bucket, preventing direct access. Configuring S3 to allow uploads from CloudFront using OAC ensures that only CloudFront can interact with the S3 bucket, adding an extra layer of security.

upvoted 9 times

George1990 **Highly Voted** 10 months ago

Correct is BD

upvoted 6 times

yangbo Most Recent 3 months ago

Selected Answer: AC

OAC do not support uploads from CloudFront to the S3 bucket.

upvoted 1 times

TicDcNess 8 months, 2 weeks ago

AWS Region for AWS Certificate Manager

To use a certificate in AWS Certificate Manager (ACM) to require HTTPS between viewers and CloudFront, make sure you request (or import) the certificate in the US East (N. Virginia) Region (us-east-1)

upvoted 2 times

rpmaws 8 months, 3 weeks ago

Selected Answer: AD

cloud front require all SSL certificate to be in us-east region regardless the origin location of the site server.

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: AD

AD looks correct

upvoted 3 times

dhewa 9 months, 3 weeks ago

Selected Answer: AD

A. Use AWS Certificate Manager (ACM) to create a public certificate in the us-east-1 Region. Use the certificate in CloudFront: CloudFront requires the certificate to be in the us-east-1 Region for custom domain names.

D. Configure Amazon S3 to allow uploads from CloudFront origin access control (OAC): This ensures secure uploads from CloudFront to the S3 bucket.

upvoted 3 times

[Removed] 9 months, 3 weeks ago

Selected Answer: AD

A. Use AWS Certificate Manager (ACM) to create a public certificate in the us-east-1 Region. Use the certificate in CloudFront.

CloudFront requires that the SSL/TLS certificate for the custom domain be created in the us-east-1 Region (N. Virginia). Even if your S3 bucket is in another region, the certificate must be in us-east-1 because CloudFront is a global service and this region is where CloudFront looks for certificates.

D. Configure Amazon S3 to allow uploads from CloudFront origin access control (OAC).

Configuring S3 to allow uploads from CloudFront using Origin Access Control (OAC) ensures that only CloudFront can interact with your S3 bucket, improving security by preventing direct access to the bucket from the public internet.

upvoted 3 times

komorebi 10 months, 1 week ago

Selected Answer: AD

Answer is AD

upvoted 4 times

JunsK1e 10 months, 1 week ago

Selected Answer: BD

BD correct answer

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #975


Topic 1

A weather forecasting company collects temperature readings from various sensors on a continuous basis. An existing data ingestion process collects the readings and aggregates the readings into larger Apache Parquet files. Then the process encrypts the files by using client-side encryption with KMS managed keys (CSE-KMS). Finally, the process writes the files to an Amazon S3 bucket with separate prefixes for each calendar day.

The company wants to run occasional SQL queries on the data to take sample moving averages for a specific calendar day.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure Amazon Athena to read the encrypted files. Run SQL queries on the data directly in Amazon S3. **Most Voted**
- B. Use Amazon S3 Select to run SQL queries on the data directly in Amazon S3.
- C. Configure Amazon Redshift to read the encrypted files. Use Redshift Spectrum and Redshift query editor v2 to run SQL queries on the data directly in Amazon S3.
- D. Configure Amazon EMR Serverless to read the encrypted files. Use Apache SparkSQL to run SQL queries on the data directly in Amazon S3.

Correct Answer: A*Community vote distribution***Comments****FlyingHawk** 4 months, 2 weeks ago**Selected Answer: A**

<https://aws.amazon.com/blogs/storage/how-to-optimize-querying-your-data-in-amazon-s3/>
upvoted 1 times

LeonSauveterre 5 months ago**Selected Answer: A**

Seems like "sample moving averages" is a typo of "simple moving averages". A simple moving average is formed by computing the average price of a security over a specific number of periods, but let's just ignore it if you don't care one bit haha

A - Athena is serverless, works perfectly with Parquet files, can handle encrypted data, and super cost-effective because you only pay for the amount of data you query.

B - Amazon S3 Select is fast for simple filtering or lightweight queries, not for complex queries like calculating moving averages.

C - Redshift with Redshift Spectrum is extremely powerful and competent to crunch heavy-duty data, so overkill for occasional SQL queries.

D - EMR is scalable and fast but is usually for heavy data pipelines, not casual SQL analysis. Like C, this is overkill.

upvoted 2 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: A

it is a simple SQL just in S3, believe it is a use case of Athena which supports client-side decryption

upvoted 1 times

Cpso 6 months, 1 week ago

Selected Answer: D

Key: CSE-KMS

A,B,C not support decrypted by CSE-KMS

upvoted 2 times

Ronanh 8 months ago

A. Athena would be overkill and more expensive for simple queries on specific files.

C. Amazon Redshift with Redshift Spectrum would be significantly more complex and costly to set up and maintain for occasional queries.

D. EMR Serverless with Apache SparkSQL would also be more complex and likely more expensive for this use case.

S3 Select provides the right balance of functionality and cost-effectiveness for the described scenario, making it the most suitable choice

upvoted 2 times

elmyth 7 months, 2 weeks ago

S3 Select is deprecated, right answer is A, clue word is "occasional", otherwise it would be Redshift.

upvoted 3 times

FlyingHawk 4 months, 2 weeks ago

Thank for sharing this, I just checked,

After careful consideration, we have made the decision to close new customer access to Amazon S3 Select and Amazon S3 Glacier Select, effective July 25, 2024. Amazon S3 Select and Amazon S3 Glacier Select existing customers can continue to use the service as usual.

upvoted 1 times

dhewa 9 months, 3 weeks ago

Selected Answer: A

A is my choice.

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: A

A sounds right

upvoted 2 times

swati1508 10 months ago

Selected Answer: A

Athena for sql

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #976

Topic 1

A company is implementing a new application on AWS. The company will run the application on multiple Amazon EC2 instances across multiple Availability Zones within multiple AWS Regions. The application will be available through the internet. Users will access the application from around the world.


The company wants to ensure that each user who accesses the application is sent to the EC2 instances that are closest to the user's location.

Which solution will meet these requirements?

- A. Implement an Amazon Route 53 geolocation routing policy. Use an internet-facing Application Load Balancer to distribute the traffic across all Availability Zones within the same Region.
- B. Implement an Amazon Route 53 geoproximity routing policy. Use an internet-facing Network Load Balancer to distribute the traffic across all Availability Zones within the same Region. **Most Voted**
- C. Implement an Amazon Route 53 multivalue answer routing policy. Use an internet-facing Application Load Balancer to distribute the traffic across all Availability Zones within the same Region.
- D. Implement an Amazon Route 53 weighted routing policy. Use an internet-facing Network Load Balancer to distribute the traffic across all Availability Zones within the same Region.

Correct Answer: B

Community vote distribution



Comments

dhewa **Highly Voted** 9 months, 3 weeks ago

Selected Answer: B

Keyword closest.

Amazon Route 53 Geoproximity Routing: This routing policy directs traffic based on the geographic location of your users and your resources, ensuring that users are routed to the closest EC2 instances.

upvoted 9 times

nebajp **Highly Voted** 10 months ago

Selected Answer: A

Correct answer is A
Geo location is based on users location
GeoProximity is based on the AWS services used by users.
upvoted 8 times

Sergantus 6 months, 3 weeks ago

The main purpose of geolocation policy is to enforce regional restrictions, provide language-specific content, or balance load across regional endpoints. It doesn't return "closest" records but only relevant (location) records.

upvoted 4 times

MortisG **Most Recent** 2 weeks, 3 days ago

Selected Answer: A

Geolocation routing ensures users are routed to their nearest region and is quite simple to set up while geoproximity routing is more suited for custom routing control, requires traffic flow and is more complex to set up.

upvoted 1 times

bora4motion 3 weeks, 2 days ago

Selected Answer: A

I am going with A. Its web, which translates to HTTP and therefore ALB.

upvoted 1 times

Salilgen 4 months, 4 weeks ago

Selected Answer: B

<https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geoproximity.html>

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: B

A - "geolocation" is simple and effective, but it might not always direct users to the physically closest Region. It's more about their geographic "label" (country/region).

B - "geoproximity" is better for ensuring actual closeness but adds a bit more complexity.

Then, I don't think ALB or NLB should be an argument here. NLBs are generally preferred for applications that require high throughput and low latency, especially if they are not using HTTP/HTTPS. NLBs operate at Layer 4 (TCP/UDP), so suitable for a wider range of applications than ALBs (Layer 7, HTTP/HTTPS). Even if the application does use HTTP/HTTPS, NLBs can still be used if you don't need the advanced features of an ALB (like content-based routing or host-based routing).

Then again, the question haven't explicitly stated the protocol so oh well, I'm gonna go with B.

upvoted 4 times

robotgeek 4 months, 1 week ago

Agree, if the question does not state "web", "ALB", or something "layer 7 specific" we should not assume it

upvoted 1 times

deacon967 5 months, 2 weeks ago

Selected Answer: A

You can argue geolocation vs geoproximity all day but why would one use a NLB for an internet based app? The is the perfect use case for an ALB.

A is the correct answer.

upvoted 1 times

LeonSauveterre 5 months ago

It's not necessary that an internet based app is bound to use ALBs for traffic balancing. If you immediately go to ALBs when a question mentions web-based apps, you're gonna get so many questions wrong.

upvoted 2 times

dragosky 5 months, 3 weeks ago

Selected Answer: B

Which AWS Route 53 routing policy should be used to route traffic based on users' location?

Geoproximity routing

Geoproximity routing lets Amazon Route 53 route traffic to your resources based on the geographic location of your users and your resources.

upvoted 1 times

8621a7c 7 months, 1 week ago

Selected Answer: A

Go with A

upvoted 2 times

b3b5fdd 7 months, 3 weeks ago

Selected Answer: B

Correct answer is B!

upvoted 4 times

Bwhizzy 7 months, 3 weeks ago

Selected Answer: A

The answer is A. routing traffic to the user region. Also, it can't be B. It says using a Network Load Balancer. you can't use a network load balancer for internet traffic

upvoted 4 times

mk168898 6 months, 3 weeks ago

TCP UDP traffic?

upvoted 2 times

kbgsgsgs 8 months, 1 week ago

Selected Answer: A

Correct answer is A

upvoted 2 times

blehbleh 8 months, 1 week ago

Selected Answer: B

It is B:

Geolocation:

Primarily focuses on the user's location, allowing you to serve different content based on their geographic region.

Geoproximity:

Considers both the user's location and the location of your resources, dynamically choosing the closest option based on network latency.

upvoted 3 times

[Removed] 9 months, 3 weeks ago

Selected Answer: B

Geolocation Routing Policy - Routes traffic based on the geographic location of the users.

Geoproximity Routing Policy - Routes traffic based on the geographic proximity of the user to AWS resources.

upvoted 4 times

komorebi 10 months ago

Selected Answer: B

Answer is B

upvoted 5 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #977

Topic 1

A financial services company plans to launch a new application on AWS to handle sensitive financial transactions. The company will deploy the application on Amazon EC2 instances. The company will use Amazon RDS for MySQL as the database. The company's security policies mandate that data must be encrypted at rest and in transit.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Configure encryption at rest for Amazon RDS for MySQL by using AWS KMS managed keys. Configure AWS Certificate Manager (ACM) SSL/TLS certificates for encryption in transit. **Most Voted**
- B. Configure encryption at rest for Amazon RDS for MySQL by using AWS KMS managed keys. Configure IPsec tunnels for encryption in transit.
- C. Implement third-party application-level data encryption before storing data in Amazon RDS for MySQL. Configure AWS Certificate Manager (ACM) SSL/TLS certificates for encryption in transit.
- D. Configure encryption at rest for Amazon RDS for MySQL by using AWS KMS managed keys. Configure a VPN connection to enable private connectivity to encrypt data in transit.

Correct Answer: A

Community vote distribution

A (100%)

Comments

JoeTromundo 8 months, 2 weeks ago

Selected Answer: A

Amazon RDS for MySQL supports encryption at rest using AWS Key Management Service (KMS) managed keys. This encryption is easy to enable during the creation of the RDS instance and requires minimal configuration. AWS KMS provides a fully managed solution for managing encryption keys, and using KMS managed keys reduces operational overhead related to key management and rotation.

Encryption in transit ensures that data transmitted between the application and the RDS database is secure. AWS Certificate Manager (ACM) can be used to provide SSL/TLS certificates, which are required to encrypt data in transit. ACM simplifies the management of SSL/TLS certificates by handling certificate renewal and deployment, reducing operational overhead.

upvoted 4 times

dhewa 9 months, 3 weeks ago

Selected Answer: A

A is my choice anyday.
upvoted 4 times

[Removed] 9 months, 3 weeks ago

Selected Answer: A

A is correct
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #978

Topic 1

A company is migrating its on-premises Oracle database to an Amazon RDS for Oracle database. The company needs to retain data for 90 days to meet regulatory requirements. The company must also be able to restore the database to a specific point in time for up to 14 days.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Create Amazon RDS automated backups. Set the retention period to 90 days.
- B. Create an Amazon RDS manual snapshot every day. Delete manual snapshots that are older than 90 days.
- C. Use the Amazon Aurora Clone feature for Oracle to create a point-in-time restore. Delete clones that are older than 90 days.
- D. Create a backup plan that has a retention period of 90 days by using AWS Backup for Amazon RDS. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

nebajp **Highly Voted** 10 months ago

Selected Answer: D

Correct Answer is D - Its fulfilling the requirement of point in time.
Automated Backup - Default retention period is 0-35 Days - so option A is wrong.
upvoted 8 times

JoeTromundo **Highly Voted** 8 months, 2 weeks ago

Selected Answer: D

A: Amazon RDS automated backups support a maximum retention period of 35 days. This option does not meet the requirement to retain backups for 90 days.
B: This approach requires manual snapshot management, including scheduling snapshots and deleting old ones. This increases operational overhead and is prone to human error.
C: This option is not applicable as Aurora Clone is a feature specific to Amazon Aurora and not available for Amazon RDS for Oracle. Additionally, it would require manual management of clones, increasing complexity.
D: AWS Backup supports point-in-time recovery for Amazon RDS, enabling you to restore the database to any specific point within the defined retention period, up to 35 days. For the requirement of 14 days, AWS Backup easily supports this capability.

upvoted 5 times

spoved **Most Recent** 8 months, 1 week ago

Selected Answer: D

<https://aws.amazon.com/getting-started/hands-on/amazon-rds-backup-restore-using-aws-backup/>

upvoted 3 times

[Removed] 9 months, 3 weeks ago

Selected Answer: D

D is right

upvoted 2 times

pujithacg8 10 months ago

Answer is D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #979

Topic 1

A company is developing a new application that uses a relational database to store user data and application configurations. The company expects the application to have steady user growth. The company expects the database usage to be variable and read-heavy, with occasional writes.

The company wants to cost-optimize the database solution. The company wants to use an AWS managed database solution that will provide the necessary performance.

Which solution will meet these requirements MOST cost-effectively?

- A. Deploy the database on Amazon RDS. Use Provisioned IOPS SSD storage to ensure consistent performance for read and write operations.
- B. Deploy the database on Amazon Aurora Serverless to automatically scale the database capacity based on actual usage to accommodate the workload. **Most Voted**
- C. Deploy the database on Amazon DynamoDB. Use on-demand capacity mode to automatically scale throughput to accommodate the workload.
- D. Deploy the database on Amazon RDS. Use magnetic storage and use read replicas to accommodate the workload.

Correct Answer: B

Community vote distribution

B (100%)

Comments

LeonSauveterre 5 months ago

Selected Answer: B

- A - You'd pay for high performance even when you don't need it
- B - You only pay for the database capacity you use, making it cost-efficient for unpredictable traffic.
- C - DynamoDB is not even a relational database.
- D - Magnetic disks have significantly slower IOPS compared to SSDs.

P.S. In the AWS context, magnetic storage refers to standard storage volumes (previously called "Magnetic" volumes) in EBS. It offers lower performance compared to modern SSD storage but is cheaper per GB.

upvoted 2 times

AMEJack 6 months, 1 week ago

Selected Answer: B

Option D is out: Aurora data is stored in the cluster volume, which is a single, virtual volume that uses solid state drives (SSDs).
So option B is correct

upvoted 2 times

[Removed] 9 months, 3 weeks ago

Selected Answer: B

B looks good

upvoted 2 times

dhewa 9 months, 3 weeks ago

Selected Answer: B

B is my choice.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #980

Topic 1

A company hosts its application on several Amazon EC2 instances inside a VPC. The company creates a dedicated Amazon S3 bucket for each customer to store their relevant information in Amazon S3.

The company wants to ensure that the application running on EC2 instances can securely access only the S3 buckets that belong to the company's AWS account.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create a gateway endpoint for Amazon S3 that is attached to the VPC. Update the IAM instance profile policy to provide access to only the specific buckets that the application needs.

B. Create a NAT gateway in a public subnet with a security group that allows access to only Amazon S3. Update the route tables to use the NAT Gateway.


C. Create a gateway endpoint for Amazon S3 that is attached to the VPC. Update the IAM instance profile policy with a Deny action and the following condition key:

```
{
  "StringNotEquals" : {
    "s3:ResourceAccount" : [ "CompanyAWSAcctNumber" ]
  }
}
```

Most Voted

D. Create a NAT Gateway in a public subnet. Update route tables to use the NAT Gateway. Assign bucket policies for all buckets with a Deny action and the following condition key:

```
{
  "StringNotEquals" : {
    "s3:ResourceAccount" : [ "CompanyAWSAcctNumber" ]
  }
}
```

Correct Answer: C*Community vote distribution***Comments**

toyaji Highly Voted 9 months ago

Selected Answer: C

B, C is not secure way because NAT gateway is for internet-facing outbound.
A is not correct because company will create dedicated bucket for each customers it means number of buckets will increase dynamically. so you cant list all on profile.

upvoted 5 times

LeonSauveterre Most Recent 5 months ago

Selected Answer: C

A - It doesn't explicitly ensure that access is limited to buckets only owned by the company. We need further adjustments to enforce this restriction (like adding a Condition key).

B & D - NAT gateways route traffic over the internet, so it's not so secure or cost-effective as S3 Gateway Endpoint.

C - Use an S3 Gateway Endpoint and restrict access using an IAM policy with a condition to deny access to any S3 resource outside the company's account. Simple, secure, and easy to understand.

upvoted 3 times

trinh_le 5 months, 1 week ago

Selected Answer: A

A. Allows option fits to minimum permission policy

C. Deny action is complexity and hard to debug

upvoted 2 times

Denise123 5 months, 1 week ago

Selected Answer: C

The requirement in the question is to ensure that the EC2 instances can securely access only the S3 buckets that belong to (are owned by) the company's AWS account.

In that context, using the "StringNotEquals": {"S3ResourceAccount":["CompanyAWSAcctNumber"]} condition key in the IAM instance profile policy is a valid approach.

This condition key restricts access to S3 buckets that are not owned by the specified AWS account number (the company's account). By setting a Deny with this condition, it effectively allows access only to the S3 buckets owned by the company's AWS account.

upvoted 3 times

Denise123 5 months, 1 week ago

Selected Answer: A

By choosing Option A, the company can leverage the simplicity and security of VPC Gateway Endpoints for Amazon S3, combined with IAM instance profile policies to grant selective access to the required S3 buckets. This solution provides the necessary isolation and access control with minimal operational overhead, making it the most efficient and scalable approach. Option C and other options have limitations or introduce additional operational overhead Option C (Creating a Gateway Endpoint and using a Deny policy with S3ResourceAccount condition) is not a valid approach. The S3ResourceAccount condition key is used to restrict access based on the AWS account that owns the S3 bucket, not the AWS account that is accessing the bucket.

upvoted 1 times

Denise123 5 months, 1 week ago

Ignore. After re-reading the question carefully, Option C is actually a valid solution.

The right answer is C and I wrote exactly that in my answer myself lol. In the question it says "BELONGS TO" (which I failed to read correctly)

The requirement in the question is to ensure that the EC2 instances can securely access only the S3 buckets that belong to (are owned by) the company's AWS account.

In that context, using the "StringNotEquals": {"S3ResourceAccount":["CompanyAWSAcctNumber"]} condition key in the IAM instance profile policy is a valid approach.

This condition key restricts access to S3 buckets that are not owned by the specified AWS account number (the company's account). By setting a Deny with this condition, it effectively allows access only to the S3 buckets owned by the company's AWS account.

upvoted 2 times

GOTJ 4 months, 1 week ago

Errare humanum est :)

upvoted 1 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: A

answer C cannot restrict ec2 from accessing s3 from other accounts.

Only by providing access to specific buckets in IAM policy on EC2 are we able to achieve this.

upvoted 1 times

AMEJack 6 months, 1 week ago

Selected Answer: A

Option C will deny the other S3 buckets but will not Allow access to the specified Buckets, thus we should have Allow rules.

Option A will allow the specified Buckets and implicit Deny the other buckets.

upvoted 2 times

ARV14 6 months, 1 week ago

You will keep updating the policy with new s3 buckets as user base grows? Operational overhead?

upvoted 1 times

youkarthik 6 months, 1 week ago

Selected Answer: A

A as per gen Als

upvoted 1 times

Bwhizzy 7 months, 3 weeks ago

Selected Answer: C

Answer is C. Just specify only the company AWS account number, rather than listing all the Buckets

upvoted 3 times

XXXXXINN 7 months, 4 weeks ago

Vote A

upvoted 2 times

siheom 8 months, 1 week ago

VOTE A

upvoted 1 times

viejito 8 months, 3 weeks ago

la respuesta correcta es la A :Los buckets son servicios globales.(o sea no están en una VPC ni Subnet), entonces no hace falta que estén en una subred publica o privada ; los Nat Gateway son para redes publicas o privadas .Entonces ahí descarta B,C y D .Cuando quieres conectar un recurso de Global Service se usa Endpoint Gateway por eso la respuesta es A .

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #981

Topic 1

A company is building a cloud-based application on AWS that will handle sensitive customer data. The application uses Amazon RDS for the database, Amazon S3 for object storage, and S3 Event Notifications that invoke AWS Lambda for serverless processing.

The company uses AWS IAM Identity Center to manage user credentials. The development, testing, and operations teams need secure access to Amazon RDS and Amazon S3 while ensuring the confidentiality of sensitive customer data. The solution must comply with the principle of least privilege.

Which solution meets these requirements with the LEAST operational overhead?

A. Use IAM roles with least privilege to grant all the teams access. Assign IAM roles to each team with customized IAM policies defining specific permission for Amazon RDS and S3 object access based on team responsibilities.

B. Enable IAM Identity Center with an Identity Center directory. Create and configure permission sets with granular access to Amazon RDS and Amazon S3. Assign all the teams to groups that have specific access with the permission sets.

Most Voted

C. Create individual IAM users for each member in all the teams with role-based permissions. Assign the IAM roles with predefined policies for RDS and S3 access to each user based on user needs. Implement IAM Access Analyzer for periodic credential evaluation.

D. Use AWS Organizations to create separate accounts for each team. Implement cross-account IAM roles with least privilege. Grant specific permission for RDS and S3 access based on team roles and responsibilities.

Correct Answer: B

Community vote distribution

B (100%)

Comments

kevin Highly Voted 9 months, 3 weeks ago

Selected Answer: B

Is there anyone who has recently passed the exam who can tell me approximately how many of the original questions are in the actual exam?

upvoted 10 times

dhewa Highly Voted 9 months, 3 weeks ago

Selected Answer: B

IAM Identity Center: This service simplifies user management by centralizing credentials and access control.

Permission Sets: You can create granular permission sets that align with the principle of least privilege, ensuring that each team has only the access they need.

Group Assignments: By assigning teams to groups with specific permission sets, you streamline access management and reduce the complexity of individual user permissions.

This approach minimizes operational overhead while maintaining secure and compliant access to sensitive customer data

upvoted 5 times

LeonSauveterre Most Recent 5 months ago

Selected Answer: B

A - Managing and customizing IAM roles for multiple teams can become more and more complex and high-maintenance, especially if team responsibilities evolve.

B - IAM Identity Center (formerly AWS SSO) allows you to manage/update access for multiple teams from one place.

C - TOO MUCH WORK.

D - Cross-account access is harder to manage and introduces additional operational overhead.

upvoted 2 times

JoeTromundo 8 months, 2 weeks ago

Selected Answer: B

Option A is good but not the best, which is option B.

upvoted 2 times

aragon_saa 9 months, 3 weeks ago

Selected Answer: B

Answer is B

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #982

Topic 1

A company has an Amazon S3 bucket that contains sensitive data files. The company has an application that runs on virtual machines in an on-premises data center. The company currently uses AWS IAM Identity Center.

The application requires temporary access to files in the S3 bucket. The company wants to grant the application secure access to the files in the S3 bucket.

Which solution will meet these requirements?

A. Create an S3 bucket policy that permits access to the bucket from the public IP address range of the company's on-premises data center.


B. Use IAM Roles Anywhere to obtain security credentials in IAM Identity Center that grant access to the S3 bucket. Configure the virtual machines to assume the role by using the AWS CLI. **Most Voted**

C. Install the AWS CLI on the virtual machine. Configure the AWS CLI with access keys from an IAM user that has access to the bucket.

D. Create an IAM user and policy that grants access to the bucket. Store the access key and secret key for the IAM user in AWS Secrets Manager. Configure the application to retrieve the access key and secret key at startup.

Correct Answer: B

Community vote distribution

**Comments**

bora4motion 3 weeks, 2 days ago

Selected Answer: D

It does not say there's a VPN to AWS or DX to AWS.

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: B

You can use AWS Identity and Access Management Roles Anywhere to obtain temporary security credentials in IAM for workloads such as servers, containers, and applications that run outside of AWS. Your workloads can use the same IAM policies

and IAM roles that you use with AWS applications to access AWS resources. Using IAM Roles Anywhere means you don't need to manage long-term credentials for workloads running outside of AWS.

<https://docs.aws.amazon.com/rolesanywhere/latest/userguide/introduction.html>

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: B

You have sensitive data in an S3 bucket, and an on-premises application needs temporary access. The solution must be secure and leverage IAM Identity Center (already in use), so now we need to grant temporary and secure credentials to the application while avoiding static credentials (like access keys).

A - IP-based access lacks identity verification, leaving the bucket vulnerable to IP spoofing.

B - This way you're not introducing static credentials or deviating from your existing centralized identity setup.

C - Manually managing credentials for multiple machines or applications is a hassle.

D - Requires additional setup (Secrets Manager integration) and introduces ongoing costs. Also, the underlying issue (static credentials) is still there.

upvoted 1 times

Bwhizzy 7 months, 3 weeks ago

Selected Answer: B

Answer is B.

AM Roles Anywhere allows on-premises servers and applications to obtain temporary AWS credentials and access AWS resources securely. This solution allows your on-premises virtual machines to use IAM roles without needing long-term credentials (like access keys). The virtual machines can assume roles and access the S3 bucket temporarily and securely.

Since the company is already using AWS IAM Identity Center, using IAM Roles Anywhere allows the company to leverage its existing Identity Center setup while following AWS best practices for security. This approach ensures the application can securely retrieve credentials without embedding static credentials into the application.

upvoted 3 times

aragon_saa 8 months, 1 week ago

Selected Answer: B

Answer is B

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #983

Topic 1

A company hosts its core network services, including directory services and DNS, in its on-premises data center. The data center is connected to the AWS Cloud using AWS Direct Connect (DX). Additional AWS accounts are planned that will require quick, cost-effective, and consistent access to these network services.

What should a solutions architect implement to meet these requirements with the LEAST amount of operational overhead?

- A. Create a DX connection in each new account. Route the network traffic to the on-premises servers.
- B. Configure VPC endpoints in the DX VPC for all required services. Route the network traffic to the on-premises servers.
- C. Create a VPN connection between each new account and the DX VPC. Route the network traffic to the on-premises servers.
- D. Configure AWS Transit Gateway between the accounts. Assign DX to the transit gateway and route network traffic to the on-premises servers. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

aragon_saa **Highly Voted** 8 months, 1 week ago

Selected Answer: D

Answer is D
upvoted 7 times

LeonSauveterre **Most Recent** 5 months ago

Selected Answer: D

New AWS accounts need consistent, quick, and cost-effective access to these services, so we can't do something for each of these accounts. We should find a way to manage these accounts fast.

A - Managing separate DX connections for every account is too much work and DX connections are expensive.

B - VPC endpoints are for accessing AWS services, not on-premises services.

C - Managing individual VPNs for multiple accounts is too much work and VPNs bring more latency.

D - Transit Gateway simplifies the network by acting as a hub for all VPCs and the on-premises network. Instead of multiple DX or VPN connections, you're using one centralized DX connection with Transit Gateway. The best part is that you can easily add new accounts or VPCs without reconfiguring DX or VPN connections.

upvoted 3 times

JA2018 5 months, 3 weeks ago

Selected Answer: D

I will choose Option D for the following reasons:

#Least operational overhead:

- Using a Transit Gateway centralizes network management, eliminating the need to create separate DX connections or VPNs for each new account, significantly reducing operational overhead.

#Scalability:

- With a Transit Gateway, you can easily add new accounts to the network without modifying existing configurations, making it ideal for scenarios with frequent account creation.

#Cost-effective:

- While creating a dedicated DX connection for each account might seem quick, it can become expensive, especially when dealing with a large number of accounts.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #984

Topic 1

A company hosts its main public web application in one AWS Region across multiple Availability Zones. The application uses an Amazon EC2 Auto Scaling group and an Application Load Balancer (ALB).

A web development team needs a cost-optimized compute solution to improve the company's ability to serve dynamic content globally to millions of customers.

Which solution will meet these requirements?

- A. Create an Amazon CloudFront distribution. Configure the existing ALB as the origin. **Most Voted**
- B. Use Amazon Route 53 to serve traffic to the ALB and EC2 instances based on the geographic location of each customer.
- C. Create an Amazon S3 bucket with public read access enabled. Migrate the web application to the S3 bucket. Configure the S3 bucket for website hosting.
- D. Use AWS Direct Connect to directly serve content from the web application to the location of each customer.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**aragon_saa** **Highly Voted** 8 months ago**Selected Answer: A**

Answer is A
upvoted 6 times

GOTJ **Most Recent** 4 months, 2 weeks ago

Cloudfront to serve dynamic content? Again?
upvoted 1 times

GOTJ 4 months, 2 weeks ago

I have to correct myself: <https://aws.amazon.com/cloudfront/dynamic-content/>
In this case, you can use cloudfront as an entry-point to AWS networking infrastructure for accessing to the web application, which is wide better than the public network infrastructure.

which is wide better than the public network infrastructure
upvoted 2 times

JA2018 6 months, 1 week ago

Selected Answer: A

CloudFront:

is a content delivery network (CDN) service from AWS, designed to distribute content globally at high speeds by caching content in edge locations close to users. This is the most cost-effective way to serve dynamic content globally, which is exactly what the scenario requires.

ALB as the origin:

By setting the existing ALB as the origin for CloudFront, the CDN will fetch content from the web application running on the EC2 instances behind the ALB, effectively caching and delivering content from geographically closer locations to users

upvoted 2 times

Cpso 6 months, 1 week ago

Selected Answer: A

"A. But I'm curious why the exam usually uses CloudFront for dynamic content. It doesn't seem reasonable."

upvoted 4 times

!!!!!!!!!!!!!!!!!!!! 6 months, 2 weeks ago

Selected Answer: A

Cloudfront == global

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #985

Topic 1

A company stores user data in AWS. The data is used continuously with peak usage during business hours. Access patterns vary, with some data not being used for months at a time. A solutions architect must choose a cost-effective solution that maintains the highest level of durability while maintaining high availability.

Which storage solution meets these requirements?

- A. Amazon S3 Standard
- B. Amazon S3 Intelligent-Tiering **Most Voted**
- C. Amazon S3 Glacier Deep Archive
- D. Amazon S3 One Zone-Infrequent Access (S3 One Zone-IA)

Correct Answer: B

Community vote distribution

B (100%)

Comments

AMEJack 6 months, 1 week ago

Selected Answer: B

Vary then intelligent tier
upvoted 2 times

Rayan999 7 months, 2 weeks ago

Selected Answer: B

B of course. finally i reached here
upvoted 4 times

viejito 8 months, 1 week ago

Respuesta correcta B . Los patrones de acceso varían = Nivel inteligente de Amazon S3 .
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #986

Topic 1

A company is testing an application that runs on an Amazon EC2 Linux instance. A single 500 GB Amazon Elastic Block Store (Amazon EBS) General Purpose SSO (gp2) volume is attached to the EC2 instance.

The company will deploy the application on multiple EC2 instances in an Auto Scaling group. All instances require access to the data that is stored in the EBS volume. The company needs a highly available and resilient solution that does not introduce significant changes to the application's code.

Which solution will meet these requirements?

- A. Provision an EC2 instance that uses NFS server software. Attach a single 500 GB gp2 EBS volume to the instance.
- B. Provision an Amazon FSx for Windows File Server file system. Configure the file system as an SMB file store within a single Availability Zone.
- C. Provision an EC2 instance with two 250 GB Provisioned IOPS SSD EBS volumes.
- D. Provision an Amazon Elastic File System (Amazon EFS) file system. Configure the file system to use General Purpose performance mode. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

GOTJ 4 months, 2 weeks ago

Selected Answer: D

In my opinion, this question is not written properly, since it doesn't clearly states that the original EBS volume contains files within a filesystem. Changing "All instances require access to the data that is stored in the EBS volume" with "All instances require access to the FILES stored in the EBS volume" would do the trick. Assuming this, "D" should be the right answer.

upvoted 1 times

JA2018 6 months, 1 week ago

Selected Answer: D

3 reasons why I will choose Option D:

#1 - Highly Available and Scalable: EFS is a managed file system designed for high availability and can automatically scale to accommodate the needs of multiple EC2 instances in an Auto Scaling group, which perfectly aligns with the company's requirement.

2 - No code changes: Since EFS is a network file system, the application doesn't need significant code changes to access the data stored on it, as it can mount the file system like a local disk.

#3 - General Purpose performance mode: This mode provides a balance between cost and performance, suitable for most application workloads.

upvoted 3 times

aragon_saa 8 months ago

Selected Answer: D

Answer is D

upvoted 4 times

aragon_saa 8 months ago

Selected Answer: D

Answer is D

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #987

Topic 1

A company recently launched a new application for its customers. The application runs on multiple Amazon EC2 instances across two Availability Zones. End users use TCP to communicate with the application.

The application must be highly available and must automatically scale as the number of users increases.

Which combination of steps will meet these requirements MOST cost-effectively? (Choose two.)

A. Add a Network Load Balancer in front of the EC2 instances. **Most Voted**

B. Configure an Auto Scaling group for the EC2 instances. **Most Voted**


C. Add an Application Load Balancer in front of the EC2 instances.

D. Manually add more EC2 instances for the application.

E. Add a Gateway Load Balancer in front of the EC2 instances.

Correct Answer: AB

Community vote distribution



Comments

blehbleh **Highly Voted** 8 months, 1 week ago

Selected Answer: AB

This is A and B. TCP best option is NLB. I know tons of you want to use chat GPT and it's going to tell you application load balancer, then ask chatgpt if an application load balancer really is the best option for TCP and it will be like "aw dang dawg, its not, you're right". Then it will switch to A and B. Because chatgpt doesn't know everything. It's a great tool but you still need to research because it doesn't have all the answers.

upvoted 8 times

classic_manda 7 months, 2 weeks ago

You are wrong! The question says, what will be the most cost effective way. ALB is much cheaper than the NLB. Thus, the best way to use NLB is when your application requires TCP and UDP connection. In this scenario, the users only communicate using TCP, which the ALB supports including variety of protocols such as, HTTP, HTTPS & SSL.

Additionally, ALB's service uptime is up to 99.995%, which falls to the requirement of this question that application must be highly available... and that nullifies your logic!

variation and that names your region.

Correct answer is C

upvoted 2 times

7dcef09 Most Recent 1 month, 3 weeks ago

Selected Answer: AB

I initially selected BC because the question mentions "application" which I immediately assumed Layer 7 communication. However, the question also mentions "End users use TCP to communicate with the application." and looking at the AWS's ELB's comparison, NLB seems like a right fit. Regardless, the question is too vague to fully understand the correct requirements. <https://aws.amazon.com/elasticloadbalancing/features/#compare>

upvoted 1 times

Dantecito 3 months, 2 weeks ago

Selected Answer: AB

TCP = NLB

highly available and automatically scale = Auto scaling group

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: AB

Even for basic TCP, ALBs are still operating at Layer 7 to some extent. They have to do more processing per connection than an NLB, which operates at Layer 4. This translates to higher resource consumption on the ALB side, even for simple TCP forwarding.

Also, both ALBs and NLBs have similar pricing structures based on

1. Hourly usage: You pay for the time the load balancer is running. This is generally comparable between the two.
2. Load Balancer Capacity Units (LCUs)

For TCP traffic, NLBs are significantly more efficient in terms of LCU consumption, especially regarding new connections and bytes processed. Even with low connection rates, the overhead of the ALB's Layer 7 processing will usually result in higher LCU consumption and therefore higher costs.

upvoted 1 times

LeonSauveterre 5 months ago

NLBs are designed to handle millions of connections per second with very low latency. Even if the application in the question doesn't require "extreme connection scaling" right now, using an NLB provides headroom for future growth without incurring significant cost increases. The cost difference between a lightly used NLB and an ALB handling the same basic TCP traffic is usually not substantial enough to justify the potential performance bottleneck and lack of scalability of the ALB for TCP. Lastly, NLB configuration for TCP is generally simpler than configuring an ALB for TCP.

upvoted 1 times

Ghoneam 5 months, 3 weeks ago

Selected Answer: BC

ALB for TCP

upvoted 1 times

siheom 7 months, 2 weeks ago

Selected Answer: BC

VOTE BC

upvoted 1 times

classic_manda 7 months, 2 weeks ago

Selected Answer: C

ALB is much cheaper than the NLB. Thus, the best way to use NLB is when your application requires TCP and UDP connection. In this scenario, the users only communicate using TCP, which the ALB supports including variety of protocols such as, HTTP, HTTPS & SSL. Additionally, ALB's service uptime is up to 99.995%, which falls to the requirement of this question that application must be highly available.

upvoted 2 times

classic_manda 7 months, 2 weeks ago

Answer is B and C

upvoted 1 times

1a0d459 8 months, 1 week ago

Selected Answer: AB

TCP - NLB

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #988

Topic 1

A company is designing the architecture for a new mobile app that uses the AWS Cloud. The company uses organizational units (OUs) in AWS Organizations to manage its accounts. The company wants to tag Amazon EC2 instances with data sensitivity by using values of sensitive and nonsensitive. IAM identities must not be able to delete a tag or create instances without a tag.

Which combination of steps will meet these requirements? (Choose two.)

A. In Organizations, create a new tag policy that specifies the data sensitivity tag key and the required values. Enforce the tag values for the EC2 instances. Attach the tag policy to the appropriate OU. **Most Voted**

B. In Organizations, create a new service control policy (SCP) that specifies the data sensitivity tag key and the required tag values. Enforce the tag values for the EC2 instances. Attach the SCP to the appropriate OU.

C. Create a tag policy to deny running instances when a tag key is not specified. Create another tag policy that prevents identities from deleting tags. Attach the tag policies to the appropriate OU.

D. Create a service control policy (SCP) to deny creating instances when a tag key is not specified. Create another SCP that prevents identities from deleting tags. Attach the SCPs to the appropriate OU. **Most Voted**

E. Create an AWS Config rule to check if EC2 instances use the data sensitivity tag and the specified values. Configure an AWS Lambda function to delete the resource if a noncompliant resource is found.

Correct Answer: AD

Community vote distribution

AD (100%)

Comments

FlyingHawk 4 months, 2 weeks ago

Selected Answer: AD

As depicted in this aws blog

<https://aws.amazon.com/blogs/mt/implement-aws-resource-tagging-strategy-using-aws-tag-policies-and-service-control-policies-scps/>

A - Tag Policy only enforces the accepted value of a tag, and not its presence. (step 1)

D - SCP Policies enforce the presence of tag. (step 2 and 3)

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: AD

B - SCPs cannot directly enforce tag values; they can only allow or deny actions based on tags.
C - Tag policies in AWS Organizations can't deny actions, such as launching untagged instances or deleting tags.
E - AWS Lambda is definitely irrelevant.
upvoted 2 times

Rayan999 7 months, 2 weeks ago

Selected Answer: AD

Ofcourse A and D. B is little bit close to the answer but its not
upvoted 3 times

jingen11 7 months, 3 weeks ago

Selected Answer: AD

A: https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps_examples_tagging.html
D: https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps_examples_tagging.html
upvoted 4 times

viejito 8 months ago

Respuesta A D .
upvoted 3 times