

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #989


Topic 1

A company runs database workloads on AWS that are the backend for the company's customer portals. The company runs a Multi-AZ database cluster on Amazon RDS for PostgreSQL.

The company needs to implement a 30-day backup retention policy. The company currently has both automated RDS backups and manual RDS backups. The company wants to maintain both types of existing RDS backups that are less than 30 days old.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure the RDS backup retention policy to 30 days for automated backups by using AWS Backup. Manually delete manual backups that are older than 30 days.
- B. Disable RDS automated backups. Delete automated backups and manual backups that are older than 30 days. Configure the RDS backup retention policy to 30 days for automated backups.
- C. Configure the RDS backup retention policy to 30 days for automated backups. Manually delete manual backups that are older than 30 days. **Most Voted**
- D. Disable RDS automated backups. Delete automated backups and manual backups that are older than 30 days automatically by using AWS CloudFormation. Configure the RDS backup retention policy to 30 days for automated backups.

Correct Answer: C*Community vote distribution***Comments****blehbleh** **Highly Voted** 8 months, 1 week ago**Selected Answer: C**

This is C, you are looking for the most cost effective solution. Again if any of you use ChatGPT it will say A because it makes the most sense for automation and less management. But it is not the most cost effective. AWS Backup costs extra money. So A, is not correct. It is C.

upvoted 11 times

XXXXXINN 7 months, 4 weeks ago

Agree!

Additionally we can configure RDS backup easily for a built-in 30 days retention, which meets the requirement and less cost than using AWS backup service

upvoted 2 times

LeonSauveterre Highly Voted 5 months ago

Selected Answer: C

A - AWS Backup is not required for configuring automated backup retention on RDS. It's native to the RDS service. So not using that will save us a fortune.

B - Disabling backups and then re-enabling a 30-day retention policy? I don't get it

C - Manual deletion will add operational tasks for sure, but in this case, it's actually reasonable and acceptable.

D - Even sounds more absurd than B, because using AWS CloudFormation to automate deletion adds extra complexity.

upvoted 5 times

jacinml Most Recent 8 months, 1 week ago

Selected Answer: D

D makes more sense, since AWS CF will keep last 30 days from now on.

upvoted 2 times

viejito 8 months, 1 week ago

.respuesta correcta A-D : Para cumplir con los requisitos de etiquetado de instancias de Amazon EC2 con sensibilidad de datos y garantizar que las identidades de IAM no puedan eliminar una etiqueta o crear instancias sin una etiqueta, la combinación de pasos más adecuada es:

A. En Organizaciones, cree una nueva política de etiqueta que especifique la clave de etiqueta de sensibilidad de datos y los valores requeridos. Aplica los valores de etiqueta para las instancias de EC2. Adjunte la política de etiquetas a la OU apropiada.

D. Cree una política de control de servicio (SCP) para denegar la creación de instancias cuando no se especifica una clave de etiqueta. Crea otro SCP que evite que las identidades eliminen las etiquetas. Adjunte los SCP a la OU apropiada.

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #990

Topic 1

A company is planning to migrate a legacy application to AWS. The application currently uses NFS to communicate to an on-premises storage solution to store application data. The application cannot be modified to use any other communication protocols other than NFS for this purpose.

Which storage solution should a solutions architect recommend for use after the migration?

- A. AWS DataSync
- B. Amazon Elastic Block Store (Amazon EBS)
- C. Amazon Elastic File System (Amazon EFS) **Most Voted**
- D. Amazon EMR File System (Amazon EMRFS)

Correct Answer: C*Community vote distribution*

C (100%)

Comments**78b9037** 6 months, 1 week ago**Selected Answer: C**

Amazon EFS is the only one which provides a fully managed NFS file system
upvoted 2 times

aragon_saa 8 months ago**Selected Answer: C**

Answer is C
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #991

Topic 1

A company uses GPS trackers to document the migration patterns of thousands of sea turtles. The trackers check every 5 minutes to see if a turtle has moved more than 100 yards (91.4 meters). If a turtle has moved, its tracker sends the new coordinates to a web application running on three Amazon EC2 instances that are in multiple Availability Zones in one AWS Region.

Recently, the web application was overwhelmed while processing an unexpected volume of tracker data. Data was lost with no way to replay the events. A solutions architect must prevent this problem from happening again and needs a solution with the least operational overhead.

What should the solutions architect do to meet these requirements?

- A. Create an Amazon S3 bucket to store the data. Configure the application to scan for new data in the bucket for processing.
- B. Create an Amazon API Gateway endpoint to handle transmitted location coordinates. Use an AWS Lambda function to process each item concurrently.
- C. Create an Amazon Simple Queue Service (Amazon SQS) queue to store the incoming data. Configure the application to poll for new messages for processing.
- D. Create an Amazon DynamoDB table to store transmitted location coordinates. Configure the application to query the table for new data for processing. Use TTL to remove data that has been processed.

Correct Answer: C*Community vote distribution*

C (100%)

Comments**Bwhizzy** 7 months, 3 weeks ago**Selected Answer: C**

The correct Answer is C. Amazon SQS is a fully managed message queuing service that allows you to decouple and scale applications by buffering the incoming data. This would ensure that data is stored in a reliable, scalable queue until the web application is ready to process it.

SQS provides high availability and fault tolerance, and guarantees that no data will be lost, as the messages remain in the queue until they are processed.

queue until they are processed.

The web application can poll the SQS queue for new messages and process them at its own pace, preventing the application from being overwhelmed by large volumes of tracker data.

Operational overhead is minimal because SQS is a fully managed service. The application only needs to poll the queue for messages, and there is no need to manage infrastructure.

upvoted 3 times

hharbiordun85 8 months ago

Answer: C

Amazon SQS provides a reliable and scalable way to decouple the components of the application. By using a queue, the incoming tracker data can be buffered, ensuring that if the web application is overwhelmed with data, the messages will still be retained in the queue until they can be processed.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #992

Topic 1

A company's software development team needs an Amazon RDS Multi-AZ cluster. The RDS cluster will serve as a backend for a desktop client that is deployed on premises. The desktop client requires direct connectivity to the RDS cluster.

The company must give the development team the ability to connect to the cluster by using the client when the team is in the office.

Which solution provides the required connectivity MOST securely?

A. Create a VPC and two public subnets. Create the RDS cluster in the public subnets. Use AWS Site-to-Site VPN with a customer gateway in the company's office.


B. Create a VPC and two private subnets. Create the RDS cluster in the private subnets. Use AWS Site-to-Site VPN with a customer gateway in the company's office. **Most Voted**

C. Create a VPC and two private subnets. Create the RDS cluster in the private subnets. Use RDS security groups to allow the company's office IP ranges to access the cluster.

D. Create a VPC and two public subnets. Create the RDS cluster in the public subnets. Create a cluster user for each developer. Use RDS security groups to allow the users to access the cluster.

Correct Answer: B

Community vote distribution



Comments

Bwhizzy **Highly Voted** 7 months, 3 weeks ago

Selected Answer: B

The Correct Answer is B.

Explanation:

VPC and Private Subnets: By placing the RDS cluster in private subnets, you ensure that the RDS cluster is not publicly accessible from the internet. This significantly improves security as the database is only accessible through secure channels, not directly from the public internet.

AWS Site-to-Site VPN: Using a Site-to-Site VPN establishes a secure, encrypted connection between the on-premises office and the AWS environment. This provides secure access to the RDS cluster without exposing it to the internet, ensuring that the developers can only access the cluster when connected to the office network.

Customer Gateway: The customer gateway is configured in the company's office to handle the VPN connection, providing secure connectivity for the desktop client to the RDS cluster when the development team is in the office.

upvoted 7 times

blehbleh Highly Voted 8 months ago

Selected Answer: B

This is B site to site von adds additional security. We are going for more secure.

upvoted 6 times

kbgsqsgs Most Recent 8 months, 1 week ago

Selected Answer: C

The goal is to limit the team to only being in the office to be in the RDS cluster, so wouldn't checking IP ranges based on the office network rather than bringing up the internet be better suited to what you really need?

upvoted 3 times

trongod05 8 months ago

But if they are in private subnets, how do they connect? Can't over public internet. And there's no connection between their office and the VPC. Needs more info I think.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #993

Topic 1


A solutions architect is creating an application that will handle batch processing of large amounts of data. The input data will be held in Amazon S3 and the output data will be stored in a different S3 bucket. For processing, the application will transfer the data over the network between multiple Amazon EC2 instances.

What should the solutions architect do to reduce the overall data transfer costs?

- A. Place all the EC2 instances in an Auto Scaling group.
- B. Place all the EC2 instances in the same AWS Region.
- C. Place all the EC2 instances in the same Availability Zone. **Most Voted**
- D. Place all the EC2 instances in private subnets in multiple Availability Zones.

Correct Answer: C

Community vote distribution



Comments

FlyingHawk 4 months, 3 weeks ago

Selected Answer: C

Data transfer within the zone is free, but cross-zone is not free:
<https://aws.amazon.com/blogs/architecture/overview-of-data-transfer-costs-for-common-architectures/>
 upvoted 1 times

ckhemani 6 months, 1 week ago

Selected Answer: C

K. After research it seems C make sense since it is saying data transfer between EC2 and data transfer between two EC2 in AZ is free but not between AZ's.

No cost within AZ: Data transfer between EC2 instances in the same AZ is free.
 Cross-AZ charges: If you transfer data between instances in different Availability Zones, you will be charged for the inter-zone data transfer.
 Region-based pricing: While data transfer within an AZ is free, data transfer across different regions may also incur charges
 upvoted 3 times

ckhemani 6 months, 1 week ago

Selected Answer: B

S3 is replicated only to 3 AZ's and you can't define AZ while creating the Bucket. If this Question would have said about Directory which is One Zone express where you can define the AZ then it makes sense to select option C.

Option B makes sense to configure EC2 in same region since EC2/S3 transfer data between same region is free too.

upvoted 2 times

78b9037 6 months, 1 week ago

Selected Answer: C

Two reasons:

1. EC2 instances in the same AZ can communicate with each other at no cost.

2. Data transfer between EC2 instances and S3 within the same AZ is free

<https://aws.amazon.com/blogs/architecture/overview-of-data-transfer-costs-for-common-architectures/>

upvoted 2 times

sOI852POL 7 months, 3 weeks ago

Selected Answer: C

C

upvoted 2 times

viejito 8 months ago

Respuesta correcta C :

upvoted 2 times

aragon_saa 8 months ago

Selected Answer: C

Answer is C

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #994

Topic 1

A company hosts a multi-tier web application that uses an Amazon Aurora MySQL DB cluster for storage. The application tier is hosted on Amazon EC2 instances. The company's IT security guidelines mandate that the database credentials be encrypted and rotated every 14 days.

What should a solutions architect do to meet this requirement with the LEAST operational effort?

- A. Create a new AWS Key Management Service (AWS KMS) encryption key. Use AWS Secrets Manager to create a new secret that uses the KMS key with the appropriate credentials. Associate the secret with the Aurora DB cluster. Configure a custom rotation period of 14 days. **Most Voted**
- B. Create two parameters in AWS Systems Manager Parameter Store: one for the user name as a string parameter and one that uses the SecureString type for the password. Select AWS Key Management Service (AWS KMS) encryption for the password parameter, and load these parameters in the application tier. Implement an AWS Lambda function that rotates the password every 14 days.
- C. Store a file that contains the credentials in an AWS Key Management Service (AWS KMS) encrypted Amazon Elastic File System (Amazon EFS) file system. Mount the EFS file system in all EC2 instances of the application tier. Restrict the access to the file on the file system so that the application can read the file and that only super users can modify the file. Implement an AWS Lambda function that rotates the key in Aurora every 14 days and writes new credentials into the file.
- D. Store a file that contains the credentials in an AWS Key Management Service (AWS KMS) encrypted Amazon S3 bucket that the application uses to load the credentials. Download the file to the application regularly to ensure that the correct credentials are used. Implement an AWS Lambda function that rotates the Aurora credentials every 14 days and uploads these credentials to the file in the S3 bucket.

Correct Answer: A

Community vote distribution

A (100%)

Comments

LeonSauveterre 5 months ago

Selected Answer: A

A - AWS KMS + Secrets Manager support custom rotation intervals, such as 14 days, with built-in integration for Aurora DB clusters

characters.

B, C, D - Manual implementation of rotation? This is really hard.

upvoted 1 times

viejito 8 months ago

respuesta correcta A . Credenciales , rotación automática =AWS Secrets Manager .

upvoted 2 times

aragon_saa 8 months ago

Selected Answer: A

Answer is A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #995

Topic 1

A streaming media company is rebuilding its infrastructure to accommodate increasing demand for video content that users consume daily.

The company needs to process terabyte-sized videos to block some content in the videos. Video processing can take up to 20 minutes.

The company needs a solution that will scale with demand and remain cost-effective.

Which solution will meet these requirements?

A. Use AWS Lambda functions to process videos. Store video metadata in Amazon DynamoDB. Store video content in Amazon S3 Intelligent-Tiering.

B. Use Amazon Elastic Container Service (Amazon ECS) and AWS Fargate to implement microservices to process videos. Store video metadata in Amazon Aurora. Store video content in Amazon S3 Intelligent-Tiering. **Most Voted**

C. Use Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer (ALB) to process videos. Store video content in Amazon S3 Standard. Use Amazon Simple Queue Service (Amazon SQS) for queuing and to decouple processing tasks.

D. Deploy a containerized video processing application on Amazon Elastic Kubernetes Service (Amazon EKS) on Amazon EC2. Store video metadata in Amazon RDS in a single Availability Zone. Store video content in Amazon S3 Glacier Deep Archive.

Correct Answer: B

Community vote distribution

B (83%)

C (17%)

Comments

kbgsqsgs **Highly Voted** 8 months, 1 week ago

Selected Answer: B

S3 Intelligent-Tiering is cost-effective for storing large amounts of video content, and since Lambda doesn't work, shouldn't we consider serverless?

upvoted 7 times

LeonSauveterre Highly Voted 5 months ago

Selected Answer: B

A - Video processing can take up to 20 minutes. So lambda is not suitable here.

B - Fargate allows running containers without managing EC2 instances, scaling automatically based on demand. Besides, S3 Intelligent-Tiering is cost-effective for storing large video files.

C - I gotta say, the SQS is a really good choice here! It's just that S3 Standard is more expensive than S3 Intelligent-Tiering for large files, and that EC2 (even with Auto Scaling) requires managing the underlying infrastructure.

D - EKS is more complex to set up but the real issue is that RDS in a single Availability Zone introduces a single point of failure, so scalability and reliability can't be met. Also, S3 Glacier Deep Archive is designed for long-term archival storage with high latency and not suitable for frequently accessed video content.

upvoted 5 times

EllenLiu Most Recent 5 months, 2 weeks ago

Selected Answer: B

it is a tough question.

terabyte-sized videos are difficult to process. we need to split the videos into chunks either with fargate or EC2.

seems DB is essential in this scenario as I have no idea how to get the videos out from s3 without metadata which should be stored into DB. So I prefer B. SQS is used for peak clipping during burst periods.

upvoted 1 times

EllenLiu 5 months, 2 weeks ago

this the whole work flow for a ideal solution:

- 1.Upload: Users upload videos to CloudFront with origin S3.
- 2.Trigger: S3 sends an event notification to SQS.
- 3.EC2 instances poll the SQS queue and download and process Video files from S3
- 4.The processed video is uploaded to another S3 bucket, metadata is saved into Aurora

upvoted 1 times

78b9037 6 months, 1 week ago

Selected Answer: B

B and C are both possible but i decided for B because:

- Fargate provides serverless containers that scale automatically
- No infrastructure management is required
- Aurora provides high availability for metadata
- S3 Intelligent-Tiering optimizes storage costs automatically
- It is not tile limited like Lambda

On the other hand C requires more management overhead it is less cost-effective than containerized solutions, and need to manage EC2 instances manually

upvoted 2 times

XXXXXINN 7 months, 4 weeks ago

This is a tough one I would say. I wish B and C can combine. but I vote for B.

For B, it is more cost efficiency focus rather than decoupling the whole process for improving overall reliability.

For C, the use of SQS is perfect solution for the downside of option B. But EC2 comes in picture which increases the cost and operational complexity.

How to pick then? lets go back to the question and see what it focuses - cost or operational complexity and stability? It looks like it leans more on focusing scalability and cost-efficiency. In that case, I would go for B because fargate provides cost-efficiency and store just metadata in DB and the rest data in S3 also provides a lower cost and improves its performance.

upvoted 2 times

XXXXXINN 7 months, 4 weeks ago

Additionally, C also introduced ALB service in the picture, that increases the cost as well.

upvoted 2 times

blehbleh 8 months ago

Selected Answer: C

C makes the most sense out of the options and given requirements.

upvoted 1 times

hbarbardun05 8 months ago

marb0r0und05 6 months ago

Answer: C

The question did not mention the application need to be containerized , i will choose C

upvoted 1 times

GOTJ 4 months, 2 weeks ago

But it mentioned "rebuilding its infrastructure", opening the door to a re-engineering process in which micro services can be applied

upvoted 1 times

blehbleh 8 months, 1 week ago

Selected Answer: C

I personally think C, I could be wrong.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #996

Topic 1

A company runs an on-premises application on a Kubernetes cluster. The company recently added millions of new customers. The company's existing on-premises infrastructure is unable to handle the large number of new customers. The company needs to migrate the on-premises application to the AWS Cloud.

The company will migrate to an Amazon Elastic Kubernetes Service (Amazon EKS) cluster. The company does not want to manage the underlying compute infrastructure for the new architecture on AWS.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use a self-managed node to supply compute capacity. Deploy the application to the new EKS cluster.
- B. Use managed node groups to supply compute capacity. Deploy the application to the new EKS cluster.
- C. Use AWS Fargate to supply compute capacity. Create a Fargate profile. Use the Fargate profile to deploy the application.
- D. Use managed node groups with Karpenter to supply compute capacity. Deploy the application to the new EKS cluster.

Most Voted**Correct Answer: C***Community vote distribution*

C (100%)

Comments**Razvan_C** **Highly Voted** 7 months, 2 weeks ago**Selected Answer: C**

"The company does not want to manage the underlying compute infrastructure for the new architecture on AWS" => Serverless
=> Fargate

upvoted 6 times

FlyingHawk **Most Recent** 4 months, 3 weeks ago**Selected Answer: C**

With Fargate, you don't have to provision, configure, or scale groups of virtual machines on your own to run containers. You also don't need to choose server types, decide when to scale your node groups, or optimize cluster packing. A Fargate profile is used to specify which Kubernetes pods should run on Fargate. By creating a Fargate profile, you can ensure that your application runs on Fargate without needing to manage node groups or EC2 instances.

upvoted 2 times

elmyth 7 months, 1 week ago

Selected Answer: C

Fargate

upvoted 3 times

viejito 8 months, 1 week ago

Respuesta correcta C = AWS Fargate = menor sobrecarga operativa. A, B y D tienen gestión de infraestructura.

upvoted 2 times

viejito 8 months, 1 week ago

Respuesta correcta C = AWS Fargate = menor sobrecarga operativa . A , B y C tienen gestión de infraestructura .

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #997

Topic 1

A company is launching a new application that requires a structured database to store user profiles, application settings, and transactional data. The database must be scalable with application traffic and must offer backups.

Which solution will meet these requirements MOST cost-effectively?

- A. Deploy a self-managed database on Amazon EC2 instances by using open source software. Use Spot Instances for cost optimization. Configure automated backups to Amazon S3.
- B. Use Amazon RDS. Use on-demand capacity mode for the database with General Purpose SSD storage. Configure automatic backups with a retention period of 7 days.
- C. Use Amazon Aurora Serverless for the database. Use serverless capacity scaling. Configure automated backups to Amazon S3. **Most Voted**
- D. Deploy a self-managed NoSQL database on Amazon EC2 instances. Use Reserved Instances for cost optimization. Configure automated backups directly to Amazon S3 Glacier Flexible Retrieval.

Correct Answer: C

Community vote distribution

C (100%)

Comments

aragon_saa **Highly Voted** 8 months ago

Selected Answer: C

Answer is C
upvoted 5 times

Itachi28 **Most Recent** 2 months ago

Selected Answer: C

answer is C
upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: C

A - Spot Instance is never a good idea for databases. These instances can be terminated with little notice, while databases are almost always persistent workloads.

B - On-Demand Capacity is suitable for steady or predictable workloads, but not as cost-efficient as Aurora Serverless for spiky or unpredictable workloads.

C - Aurora Serverless is a fully managed relational database that scales automatically with traffic. Automated Backups is also integrated with Aurora. Good!

D - We need a structured database, not NoSQL database.

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #998


Topic 1

A company runs its legacy web application on AWS. The web application server runs on an Amazon EC2 instance in the public subnet of a VPC. The web application server collects images from customers and stores the image files in a locally attached Amazon Elastic Block Store (Amazon EBS) volume. The image files are uploaded every night to an Amazon S3 bucket for backup.

A solutions architect discovers that the image files are being uploaded to Amazon S3 through the public endpoint. The solutions architect needs to ensure that traffic to Amazon S3 does not use the public endpoint.

Which solution will meet these requirements?

- A. Create a gateway VPC endpoint for the S3 bucket that has the necessary permissions for the VPC. Configure the subnet route table to use the gateway VPC endpoint. **Most Voted**
- B. Move the S3 bucket inside the VPC. Configure the subnet route table to access the S3 bucket through private IP addresses.
- C. Create an Amazon S3 access point for the Amazon EC2 instance inside the VPC. Configure the web application to upload by using the Amazon S3 access point.
- D. Configure an AWS Direct Connect connection between the VPC that has the Amazon EC2 instance and Amazon S3 to provide a dedicated network path.

Correct Answer: A*Community vote distribution***Comments****FlyingHawk** 4 months, 3 weeks ago**Selected Answer: A**

A gateway VPC endpoint allows private connectivity between your VPC and Amazon S3. It ensures that traffic between the VPC and S3 stays within the AWS network and does not traverse the public internet.

B - S3 buckets cannot be "moved inside the VPC." S3 is a regional service and does not reside within a VPC. This option is not valid.

C - S3 access points are used to simplify access control for S3 buckets, but they do not provide private connectivity. Traffic to S3 access points still uses the public endpoint unless combined with a gateway VPC endpoint.

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

D- Direct Connect provides a dedicated network connection between on-premises environments and AWS. It is not necessary for connecting resources within AWS (e.g., EC2 instances in a VPC to S3). This option is overkill and adds unnecessary cost and complexity.

upvoted 1 times

LeonSauveterre 5 months ago

Selected Answer: A

A - Gateway VPC Endpoint provides private connectivity between resources in a VPC and AWS services like S3 without requiring an internet gateway, NAT gateway, or public IP address.

B - This is just bull. Amazon S3 is a global service that cannot "move" into a VPC.

C - This doesn't ensure private traffic between the EC2 instance and S3 at all. You still need VPC Endpoint.

D - This may actually work but more expensive and complex compared to option A.

upvoted 3 times

LeonSauveterre 5 months ago

About connections:

1. VPN Connection: For secure, encrypted traffic between on-premises and AWS over the public internet.

2. Direct Connect: For dedicated private network connectivity between on-premises and AWS (like data centers with large workloads).

3. PrivateLink: For secure private connectivity to AWS services or third-party services (like Salesforce) within AWS.

4. VPC Peering: For direct connectivity between two AWS VPCs (like multi-account or region scenarios).

upvoted 3 times

LeonSauveterre 5 months ago

5. VPC Endpoint (Gateway): Provides private connectivity between VPC and services like S3 and DynamoDB.

6. VPC Endpoint (Interface): Connects to services over private IP using an ENI in your VPC. Suitable for services like SNS, SQS, or KMS.

7. S3 Access Point: Simplifies access to S3 buckets with distinct policies for specific use cases.

8. Internet Gateway: Allows resources in a VPC to connect to the internet.

9. NAT Gateway/Instance: Allows private instances to connect to the internet or public AWS services. Aka, enables outbound internet access for instances in private subnets.

upvoted 4 times

dfgdsfgfdgreg 4 months, 2 weeks ago

Correct. The S3 access point provides access management, but not a private connection. A VPC endpoint allows an instance to directly upload to the S3 whereas the S3 access point still requires the EC2 to travel over the internet.

upvoted 2 times

JA2018 5 months, 3 weeks ago

Selected Answer: C

I will choose Option C for the following reasons:

Private access:

S3 access points allow you to create a private endpoint within your VPC that can be used to access the S3 bucket without going through the public endpoint.

Security best practice:

Using an S3 access point is considered a secure way to manage access to your S3 buckets from within your VPC.

Configuration simplicity:

You only need to configure the web application to use the S3 access point, making it a relatively straightforward implementation.

upvoted 3 times

aragon_saa 8 months ago

Selected Answer: A

Answer is A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #999

Topic 1

A company is creating a prototype of an ecommerce website on AWS. The website consists of an Application Load Balancer, an Auto Scaling group of Amazon EC2 instances for web servers, and an Amazon RDS for MySQL DB instance that runs with the Single-AZ configuration.

The website is slow to respond during searches of the product catalog. The product catalog is a group of tables in the MySQL database that the company does not update frequently. A solutions architect has determined that the CPU utilization on the DB instance is high when product catalog searches occur.

What should the solutions architect recommend to improve the performance of the website during searches of the product catalog?

A. Migrate the product catalog to an Amazon Redshift database. Use the COPY command to load the product catalog tables.

B. Implement an Amazon ElastiCache for Redis cluster to cache the product catalog. Use lazy loading to populate the cache. **Most Voted**

C. Add an additional scaling policy to the Auto Scaling group to launch additional EC2 instances when database response is slow.

D. Turn on the Multi-AZ configuration for the DB instance. Configure the EC2 instances to throttle the product catalog queries that are sent to the database.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**aragon_saa** **Highly Voted** 8 months, 1 week ago**Selected Answer: B**

Answer is B
upvoted 7 times

LeonSauveterre **Most Recent** 5 months ago

Selected Answer: B

A - For transactional or low-latency queries like product catalog searches in an ecommerce website, Redshift is not suitable because it's a data warehouse solution optimized for analytical queries.

B - ElastiCache for Redis is a highly performant, in-memory caching service that can significantly reduce database load by caching frequent queries.

C - Adding more web servers can't help alleviate the load on the database. The database remains the bottleneck for product catalog queries.

D - If you turn on throttle queries, customers may gradually ask for a refund :) Let's get rid of the bad things in database instead of user experience.

BTW, what Lazy Loading means: Data is added to the cache only when requested. If not found in the cache, it is fetched from the database and added to the cache for subsequent requests.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1000

Topic 1


A company currently stores 5 TB of data in on-premises block storage systems. The company's current storage solution provides limited space for additional data. The company runs applications on premises that must be able to retrieve frequently accessed data with low latency. The company requires a cloud-based storage solution.

Which solution will meet these requirements with the MOST operational efficiency?

- A. Use Amazon S3 File Gateway. Integrate S3 File Gateway with the on-premises applications to store and directly retrieve files by using the SMB file system.
- B. Use an AWS Storage Gateway Volume Gateway with cached volumes as iSCSI targets. **Most Voted**
- C. Use an AWS Storage Gateway Volume Gateway with stored volumes as iSCSI targets.
- D. Use an AWS Storage Gateway Tape Gateway. Integrate Tape Gateway with the on-premises applications to store virtual tapes in Amazon S3.

Correct Answer: B

Community vote distribution



Comments

sOI852POL **Highly Voted** 7 months, 3 weeks ago

Selected Answer: B

Question says, "Apps must be able to retrieve frequently accessed data with low latency" so we go with cached volumes. We'd have chosen the stored volumes if question was about low-latency access to the entire dataset.

upvoted 11 times

FlyingHawk **Most Recent** 4 months, 3 weeks ago

Selected Answer: B

B - Cached volumes store only the most frequently accessed data on-premises, while the full dataset is stored in Amazon S3. This ensures low-latency access to frequently accessed data while leveraging the scalability and durability of S3 for the full dataset.

C - Stored volumes store the entire dataset on-premises and asynchronously back it up to AWS. This option does not address the company's requirement for limited on-premises storage capacity.

upvoted 2 times

FlyingHawk 4 months, 3 weeks ago

A - S3 File Gateway provides file-based storage (SMB or NFS) and is not suitable for block storage requirements.
D - Tape Gateway is designed for backup and archival use cases, not for low-latency access to frequently accessed data. It uses virtual tapes stored in S3 and is not suitable for block storage or real-time application access.
upvoted 2 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: C

choose store volume mode
upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Stored volumes store the entire dataset on-premises and asynchronously back it up to AWS. This option does not address the company's requirement for limited on-premises storage capacity, as it requires storing the full dataset locally.
upvoted 1 times

8621a7c 7 months, 3 weeks ago

Selected Answer: C

Block storage use volume gateway, low latency use store volumes.
upvoted 2 times

viejito 8 months, 1 week ago

Respuesta correcta B . Baja latencia = Elastic Cache . A , C y D no cumplen un acceso con baja latencia (descartados) .
upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1001

Topic 1

A company operates a food delivery service. Because of recent growth, the company's order processing system is experiencing scaling problems during peak traffic hours. The current architecture includes Amazon EC2 instances in an Auto Scaling group that collect orders from an application. A second group of EC2 instances in an Auto Scaling group fulfills the orders.

The order collection process occurs quickly, but the order fulfillment process can take longer. Data must not be lost because of a scaling event.

A solutions architect must ensure that the order collection process and the order fulfillment process can both scale adequately during peak traffic hours.

Which solution will meet these requirements?

- A. Use Amazon CloudWatch to monitor the CPUUtilization metric for each instance in both Auto Scaling groups. Configure each Auto Scaling group's minimum capacity to meet its peak workload value.
- B. Use Amazon CloudWatch to monitor the CPUUtilization metric for each instance in both Auto Scaling groups. Configure a CloudWatch alarm to invoke an Amazon Simple Notification Service (Amazon SNS) topic to create additional Auto Scaling groups on demand.
- C. Provision two Amazon Simple Queue Service (Amazon SQS) queues. Use one SQS queue for order collection. Use the second SQS queue for order fulfillment. Configure the EC2 instances to poll their respective queues. Scale the Auto Scaling groups based on notifications that the queues send.
- D. Provision two Amazon Simple Queue Service (Amazon SQS) queues. Use one SQS queue for order collection. Use the second SQS queue for order fulfillment. Configure the EC2 instances to poll their respective queues. Scale the Auto Scaling groups based on the number of messages in each queue. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

sOI852POL 7 months, 3 weeks ago

Selected Answer: D

SQS to ensure data is not lost because of a scaling event, scale based on number of messages on the queue.

upvoted 4 times

8621a7c 7 months, 3 weeks ago

Selected Answer: D

Go with D

upvoted 2 times

aragon_saa 8 months ago

Selected Answer: D

Answer is D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1002

Topic 1

An online gaming company is transitioning user data storage to Amazon DynamoDB to support the company's growing user base. The current architecture includes DynamoDB tables that contain user profiles, achievements, and in-game transactions.


The company needs to design a robust, continuously available, and resilient DynamoDB architecture to maintain a seamless gaming experience for users.

Which solution will meet these requirements MOST cost-effectively?

- A. Create DynamoDB tables in a single AWS Region. Use on-demand capacity mode. Use global tables to replicate data across multiple Regions.
- B. Use DynamoDB Accelerator (DAX) to cache frequently accessed data. Deploy tables in a single AWS Region and enable auto scaling. Configure Cross-Region Replication manually to additional Regions.
- C. Create DynamoDB tables in multiple AWS Regions. Use on-demand capacity mode. Use DynamoDB Streams for Cross-Region Replication between Regions.
- D. Use DynamoDB global tables for automatic multi-Region replication. Deploy tables in multiple AWS Regions. Use provisioned capacity mode. Enable auto scaling. **Most Voted**

Correct Answer: D

Community vote distribution



Comments

trinh_le 6 months ago

Selected Answer: D

Following keywords

robust: provisioned capacity reduces cost

continuously available: global tables

resilient: multiple regions enhances available

upvoted 2 times

Cpso 6 months, 1 week ago

Selected Answer: D

c & d can do 2 ways rep. But C need user logic to handle conflict ,so it not robust
upvoted 2 times

78b9037 6 months, 2 weeks ago

Selected Answer: D

<https://docs.aws.amazon.com/wellarchitected/latest/serverless-applications-lens/capacity.html>
upvoted 2 times

striker89 7 months ago

Selected Answer: C

DynamoDB Streams for replication are less expensive than Global tables.
On-demand capacity mode can be less expensive than provisioned mode.
Multi region deployment ensure high availability.
upvoted 1 times

Bwhizzy 7 months, 3 weeks ago

Selected Answer: D

D. Use DynamoDB global tables for automatic multi-Region replication. Deploy tables in multiple AWS Regions. Use provisioned capacity mode. Enable auto scaling.

Explanation:

DynamoDB global tables automatically replicate data across multiple Regions, ensuring that the data is available and consistent across all Regions. This provides resilience and high availability by allowing users in different geographical locations to access data from the closest Region.

Provisioned capacity mode allows you to pre-allocate read and write capacity units, which can result in cost savings over on-demand capacity mode if the traffic is predictable. Additionally, auto scaling can be enabled to dynamically adjust the capacity based on the actual traffic, ensuring that you only pay for the capacity that you need.

Multi-Region deployment improves the resilience of the system. If a failure occurs in one Region, another Region can seamlessly take over, ensuring an uninterrupted gaming experience.

upvoted 4 times

GOTJ 4 months, 3 weeks ago

"if the traffic is predictable". Good point, but how do you know this from the question?

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1003

Topic 1

A company runs its media rendering application on premises. The company wants to reduce storage costs and has moved all data to Amazon S3. The on-premises rendering application needs low-latency access to storage.

The company needs to design a storage solution for the application. The storage solution must maintain the desired application performance.

Which storage solution will meet these requirements in the MOST cost-effective way?

- A. Use Mountpoint for Amazon S3 to access the data in Amazon S3 for the on-premises application.
- B. Configure an Amazon S3 File Gateway to provide storage for the on-premises application. **Most Voted**
- C. Copy the data from Amazon S3 to Amazon FSx for Windows File Server. Configure an Amazon FSx File Gateway to provide storage for the on-premises application.
- D. Configure an on-premises file server. Use the Amazon S3 API to connect to S3 storage. Configure the application to access the storage from the on-premises file server.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**LeonSauveterre** 5 months ago**Selected Answer: B**

A - Mountpoint is not designed for on-premises environments or low-latency access; it works well for cloud-based applications.
B - Amazon S3 File Gateway provides low-latency, on-premises access to data stored in Amazon S3 through a file-based interface. It caches frequently accessed data locally to reduce latency.

C - Copying data to FSx increases costs, as FSx storage is more expensive than S3. Let alone unnecessary complexity.

D - Direct interaction with S3 introduces higher latency for on-premises applications and lacks local caching.

upvoted 4 times

Americanman 6 months, 2 weeks ago

Amazon S3 File Gateway enables your existing file-based applications, devices, and workflows to use Amazon S3, without modification. Amazon S3 File Gateway securely and durably stores both file contents and metadata as objects, while providing your on-premises applications low-latency access to cached data

your on-premises applications low-latency access to cached data.
upvoted 2 times

Bwhizzy 7 months, 3 weeks ago

Selected Answer: B

B. Configure an Amazon S3 File Gateway to provide storage for the on-premises application.

Explanation:

Amazon S3 File Gateway provides a way for on-premises applications to access objects stored in Amazon S3 as files. It caches frequently accessed data locally, which ensures low-latency access to the data. This is crucial for maintaining the performance of the rendering application.

By keeping the data in Amazon S3, the company benefits from lower storage costs compared to using other storage services like Amazon FSx, while still providing the necessary performance for the on-premises application through the local caching capabilities of the File Gateway.

The File Gateway seamlessly integrates with Amazon S3, allowing the application to access data using standard file protocols like NFS or SMB, which simplifies the setup.

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions


Question #1004

Topic 1

A company hosts its enterprise resource planning (ERP) system in the us-east-1 Region. The system runs on Amazon EC2 instances. Customers use a public API that is hosted on the EC2 instances to exchange information with the ERP system. International customers report slow API response times from their data centers.

Which solution will improve response times for the international customers MOST cost-effectively?

- A. Create an AWS Direct Connect connection that has a public virtual interface (VIF) to provide connectivity from each customer's data center to us-east-1. Route customer API requests by using a Direct Connect gateway to the ERP system API.
- B. Set up an Amazon CloudFront distribution in front of the API. Configure the CachingOptimized managed cache policy to provide improved cache efficiency.
- C. Set up AWS Global Accelerator. Configure listeners for the necessary ports. Configure endpoint groups for the appropriate Regions to distribute traffic. Create an endpoint in the group for the API. **Most Voted**
- D. Use AWS Site-to-Site VPN to establish dedicated VPN tunnels between Regions and customer networks. Route traffic to the API over the VPN connections.

Correct Answer: C*Community vote distribution***Comments****TewatiaAmit** **Highly Voted** 7 months, 3 weeks ago**Selected Answer: C**

CloudFront can reduce response times by caching API responses, but if the API is dynamic and not cacheable, it may not be as effective. Global Accelerator is better for improving latency when caching is not an option.

upvoted 11 times

Sergantus 6 months, 3 weeks ago

While dynamic content typically has low caching potential, CloudFront reduces latency by routing requests to the nearest edge location. There is also TCP Connection Reuse, which is also beneficial for low latency.

upvoted 2 times

8621a7c Highly Voted 7 months, 3 weeks ago

Selected Answer: B

CloudFront for Dynamic content (such as API acceleration and dynamic site delivery)

upvoted 10 times

tch Most Recent 2 months, 4 weeks ago

Selected Answer: B

How is AWS Global Accelerator different from Amazon CloudFront?

AWS Global Accelerator and Amazon CloudFront are separate services that use the AWS global network and its edge locations around the world. CloudFront improves performance for both cacheable content (such as images and videos) and dynamic content (such as API acceleration and dynamic site delivery). Global Accelerator improves performance for a wide range of applications over TCP or UDP by proxying packets at the edge to applications running in one or more AWS Regions. Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover. Both services integrate with AWS Shield for DDoS protection

upvoted 1 times

tch 2 months, 4 weeks ago

MOST cost-effectively.. this is another reason to use CloudFont...

upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: B

Given that the ERP system is hosted in us-east-1 and the API is likely HTTP/HTTPS-based, Option B (Amazon CloudFront) is the best solution. CloudFront's global edge network and caching capabilities will significantly improve response times for international customers while reducing the load on the backend EC2 instances.

upvoted 1 times

joanna91 4 months, 3 weeks ago

Selected Answer: C

C. AWS global accelerator is more suitable for cost-effective solutions to improve API response time for int'l users by using AWS global network to optimize traffic routing. Whereas, D. AWS Cloudfront is less effective for dynamic, real-time API interactions like this ERP system.

upvoted 2 times

LeonSauveterre 5 months ago

Selected Answer: C

A - Direct Connect requires dedicated connections for each customer, making it expensive and impractical for a public API.
B - CloudFront improves performance but caching efficiency depends on the API use case. Dynamic & personalized API responses may not benefit significantly from caching. CloudFront is best suited for static content or APIs with predictable responses.

C - Global Accelerator provides a global static IP address and uses the AWS global network to route traffic optimally, reducing latency for geographically distant users. Unlike CloudFront, it works well in this case because it doesn't rely on caching contents.
D - VPNs are not ideal for public APIs because they require specific setup for each customer, so high operational overhead.

upvoted 2 times

GOTJ 5 months ago

Selected Answer: C

I think that the following quote "to exchange information with the ERP system" suggest that the clients use the api not only to get, but to add and/or update ERP data. If this is the case, using a Cloudfront distribution for caching should be useless as the add/update operations must be performed by the ERP system itself.

Option "C" approach is different: it's "forcing" users to access API ERP via AWS routing optimised networks by configuring the current API as an endpoint for the globally distributed endpoints groups (no need to duplicate current infrastructure or setup replication)

upvoted 3 times

Anyio 5 months, 1 week ago

Selected Answer: C

The correct answer is C.

Explanation:

Option C: Correct. AWS Global Accelerator provides access via the AWS global network, reducing latency and improving performance for globally distributed users. By using Global Accelerator, you can direct global traffic to optimal AWS endpoints, offering lower latency and better connectivity compared to standard internet routing. It achieves this without modifying the API infrastructure, making it a highly cost-effective solution for improving API response times for international users.

upvoted 1 times

tch 5 months, 2 weeks ago

Selected Answer: B

MOST cost-effectively -
(ERP) system in the us-east-1 Region (one region only)
will be B

AWS Global Accelerator use case for one or multiple AWS regions!

upvoted 1 times

dragossky 5 months, 3 weeks ago

Selected Answer: C

C. Set up AWS Global Accelerator. Configure listeners for the necessary ports. Configure endpoint groups for the appropriate Regions to distribute traffic. Create an endpoint in the group for the API.

upvoted 3 times

DhirajBansal 6 months ago

Selected Answer: B

Cloudfront used for API acceleration and Dynamic content delivery

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1005

Topic 1

A company tracks customer satisfaction by using surveys that the company hosts on its website. The surveys sometimes reach thousands of customers every hour. Survey results are currently sent in email messages to the company so company employees can manually review results and assess customer sentiment.

The company wants to automate the customer survey process. Survey results must be available for the previous 12 months.

Which solution will meet these requirements in the MOST scalable way?

A. Send the survey results data to an Amazon API Gateway endpoint that is connected to an Amazon Simple Queue Service (Amazon SQS) queue. Create an AWS Lambda function to poll the SQS queue, call Amazon Comprehend for sentiment analysis, and save the results to an Amazon DynamoDB table. Set the TTL for all records to 365 days in the future.

Most Voted

B. Send the survey results data to an API that is running on an Amazon EC2 instance. Configure the API to store the survey results as a new record in an Amazon DynamoDB table, call Amazon Comprehend for sentiment analysis, and save the results in a second DynamoDB table. Set the TTL for all records to 365 days in the future.

C. Write the survey results data to an Amazon S3 bucket. Use S3 Event Notifications to invoke an AWS Lambda function to read the data and call Amazon Rekognition for sentiment analysis. Store the sentiment analysis results in a second S3 bucket. Use S3 lifecycle policies on each bucket to expire objects after 365 days.

D. Send the survey results data to an Amazon API Gateway endpoint that is connected to an Amazon Simple Queue Service (Amazon SQS) queue. Configure the SQS queue to invoke an AWS Lambda function that calls Amazon Lex for sentiment analysis and saves the results to an Amazon DynamoDB table. Set the TTL for all records to 365 days in the future.

Correct Answer: A

Community vote distribution

A (100%)

Comments

EllenLiu 5 months, 2 weeks ago

Selected Answer: A

<https://docs.aws.amazon.com/comprehend/latest/dg/what-is.html>

Upvoted 1 times

upvoted 1 times

trinh_le 6 months ago

Selected Answer: A

Following keywords
reach thousands of customers every hour: use SQS
Automate handle survey: aws comprehen
upvoted 1 times

siheom 6 months, 1 week ago

Selected Answer: A

VOTE A

upvoted 2 times

Changwha 6 months, 2 weeks ago

A. This solution is the most scalable, cost-effective, and tailored to the requirements. It uses Amazon Comprehend for sentiment analysis, leverages serverless architecture, and retains data for 12 months using DynamoDB's TTL feature.
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1006


Topic 1

A company uses AWS Systems Manager for routine management and patching of Amazon EC2 instances. The EC2 instances are in an IP address type target group behind an Application Load Balancer (ALB).

New security protocols require the company to remove EC2 instances from service during a patch. When the company attempts to follow the security protocol during the next patch, the company receives errors during the patching window.

Which combination of solutions will resolve the errors? (Choose two.)

- A. Change the target type of the target group from IP address type to instance type.
- B. Continue to use the existing Systems Manager document without changes because it is already optimized to handle instances that are in an IP address type target group behind an ALB.
- C. Implement the `AWSEC2-PatchLoadBalancerInstance` Systems Manager Automation document to manage the patching process. **Most Voted**
- D. Use Systems Manager Maintenance Windows to automatically remove the instances from service to patch the instances. **Most Voted**
- E. Configure Systems Manager State Manager to remove the instances from service and manage the patching schedule. Use ALB health checks to re-route traffic.

Correct Answer: CD*Community vote distribution***Comments****LeonSauveterre** **Highly Voted** 5 months ago**Selected Answer:** CD

YOU SHOULD KNOW: The errors occur because the company is trying to patch instances that are still actively serving traffic.

A - Unnecessary when existing tools can handle the issue.

B - The existing document is not optimized for instances in an IP address type target group, so errors would persist.

C - This document automates the process of deregistering instances from the ALB during patching and re-registering them afterward, ensuring no traffic is routed to the instances while they are being patched. It directly addresses the issue and adheres to best practices.

D - Maintenance Windows allow you to define a schedule for patching and ensure instances are temporarily removed from service during the process.

E - This alone can't de-register and re-register instances.

upvoted 5 times

JoeAWS Most Recent 3 months, 1 week ago

Selected Answer: AC

The document deregisters instances from target group. Since they are using IP address type, it won't prevent traffic from reaching the server. Need to work choice A into the answer.

upvoted 1 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: CD

<https://docs.aws.amazon.com/systems-manager/latest/userguide/state-manager-vs-maintenance-windows.html>

upvoted 1 times

bujuman 6 months, 3 weeks ago

Selected Answer: CD

<https://docs.aws.amazon.com/systems-manager-automation-runbooks/latest/userguide/automation-awsec2-patch-load-balancer-instance.html>

upvoted 3 times

aragon_saa 6 months, 3 weeks ago

Selected Answer: CD

Answer is CD

upvoted 3 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1007

Topic 1

A medical company wants to perform transformations on a large amount of clinical trial data that comes from several customers. The company must extract the data from a relational database that contains the customer data. Then the company will transform the data by using a series of complex rules. The company will load the data to Amazon S3 when the transformations are complete.


All data must be encrypted where it is processed before the company stores the data in Amazon S3. All data must be encrypted by using customer-specific keys.

Which solution will meet these requirements with the LEAST amount of operational effort?

- A. Create one AWS Glue job for each customer. Attach a security configuration to each job that uses server-side encryption with Amazon S3 managed keys (SSE-S3) to encrypt the data.
- B. Create one Amazon EMR cluster for each customer. Attach a security configuration to each cluster that uses client-side encryption with a custom client-side root key (CSE-Custom) to encrypt the data.
- C. Create one AWS Glue job for each customer. Attach a security configuration to each job that uses client-side encryption with AWS KMS managed keys (CSE-KMS) to encrypt the data. **Most Voted**
- D. Create one Amazon EMR cluster for each customer. Attach a security configuration to each cluster that uses server-side encryption with AWS KMS keys (SSE-KMS) to encrypt the data.

Correct Answer: C

Community vote distribution



Comments

bora4motion 1 month ago

Selected Answer: D

I am tempted to go with C but does it support CSE-KMS ?
upvoted 1 times

FlyingHawk 4 months, 3 weeks ago

Selected Answer: C

Option C is the best solution because it uses AWS Glue for serverless ETL processing and CSE-KMS for end-to-end encryption using customer-specific keys

upvoted 3 times

sOI852POL 7 months, 3 weeks ago

Selected Answer: C

AWS Glue for ETL, then AWS KMS (CSE-KMS)

upvoted 4 times

TewatiaAmit 7 months, 3 weeks ago

Selected Answer: C

AWS Glue with client-side encryption using AWS KMS (CSE-KMS) provides the required pre-processing encryption with minimal operational effort.

upvoted 2 times

blehbleh 8 months ago

Selected Answer: C

It's C gotta use glue and since it's before the company stores the data in Amazon S3 gotta be client side.

upvoted 4 times

blehbleh 8 months, 1 week ago

Selected Answer: C

I say C, aws glue reduces the operational management uses server side encryption and ins which allows for user specific keys.

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1008

Topic 1

A company hosts a website analytics application on a single Amazon EC2 On-Demand Instance. The analytics application is highly resilient and is designed to run in stateless mode.

The company notices that the application is showing signs of performance degradation during busy times and is presenting 5xx errors. The company needs to make the application scale seamlessly.

Which solution will meet these requirements MOST cost-effectively?

- A. Create an Amazon Machine Image (AMI) of the web application. Use the AMI to launch a second EC2 On-Demand Instance. Use an Application Load Balancer to distribute the load across the two EC2 instances.
- B. Create an Amazon Machine Image (AMI) of the web application. Use the AMI to launch a second EC2 On-Demand Instance. Use Amazon Route 53 weighted routing to distribute the load across the two EC2 instances.
- C. Create an AWS Lambda function to stop the EC2 instance and change the instance type. Create an Amazon CloudWatch alarm to invoke the Lambda function when CPU utilization is more than 75%.
- D. Create an Amazon Machine Image (AMI) of the web application. Apply the AMI to a launch template. Create an Auto Scaling group that includes the launch template. Configure the launch template to use a Spot Fleet. Attach an Application Load Balancer to the Auto Scaling group. **Most Voted**

Correct Answer: D

Community vote distribution

D (100%)

Comments

LeonSauveterre 4 months, 4 weeks ago

Selected Answer: D

A - This can't scale automatically during peak times.

B - Weighted routing is not ideal for seamless scaling because it requires manual intervention to adjust weights or add new instances.

C - It introduces downtime during instance type changes, let alone scaling.

D - (1) ASG dynamically adjusts the number of instances based on demand. (2) Spot Instances reduces costs significantly compared to On-Demand Instances. (3) ALB + ASG ensures efficient traffic routing to healthy instances.

upvoted 3 times

JA2018 6 months, 1 week ago

Selected Answer: D

3 reasons why I will choose Option D:

1. Scalability: Auto Scaling groups are designed to automatically scale up or down the number of instances based on demand, which is ideal for handling fluctuating traffic during busy periods.
 2. Cost-effectiveness: Using a Spot Fleet allows you to leverage unused EC2 capacity at significantly lower costs compared to On-Demand instances.
 3. Stateless Application: Since the application is stateless, it can easily be scaled horizontally with multiple instances behind an Application Load Balancer without any data consistency issues
- upvoted 2 times

Razvan_C 7 months, 1 week ago

Selected Answer: D

With auto scalling group it can also scale down on low demand (costs saved)

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1009

Topic 1

A company runs an environment where data is stored in an Amazon S3 bucket. The objects are accessed frequently throughout the day. The company has strict data encryption requirements for data that is stored in the S3 bucket. The company currently uses AWS Key Management Service (AWS KMS) for encryption.


The company wants to optimize costs associated with encrypting S3 objects without making additional calls to AWS KMS.

Which solution will meet these requirements?

- A. Use server-side encryption with Amazon S3 managed keys (SSE-S3).
- B. Use an S3 Bucket Key for server-side encryption with AWS KMS keys (SSE-KMS) on the new objects. **Most Voted**
- C. Use client-side encryption with AWS KMS customer managed keys.
- D. Use server-side encryption with customer-provided keys (SSE-C) stored in AWS KMS.

Correct Answer: B

Community vote distribution



Comments

8621a7c **Highly Voted** 7 months, 3 weeks ago

Selected Answer: B

reduce cost <https://docs.aws.amazon.com/AmazonS3/latest/userguide/bucket-key.html>

upvoted 8 times

yangbo **Most Recent** 3 months ago

Selected Answer: A

without making additional calls to AWS KMS

upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

Selected Answer: B

When you configure your bucket to use an S3 Bucket Key for SSE-KMS, AWS generates a short-lived bucket-level key from AWS KMS, then temporarily keeps it in S3. This bucket-level key will create data keys for new objects during its lifecycle. S3 Bucket

KMS, then temporarily keeps it in S3. This bucket-level key will create data keys for new objects during its lifecycle. S3 Bucket Keys are used for a limited time period within Amazon S3, reducing the need for S3 to make requests to AWS KMS to complete encryption operations. This reduces traffic from S3 to AWS KMS, allowing you to access AWS KMS-encrypted objects in Amazon S3 at a fraction of the previous cost.

upvoted 2 times

FlyingHawk 4 months, 2 weeks ago

Selected Answer: B

Using S3 Bucket Keys allows you to save on AWS KMS request costs by decreasing your requests to AWS KMS for Encrypt, GenerateDataKey, and Decrypt operations through the use of a bucket-level key. By design, subsequent requests that take advantage of this bucket-level key do not result in AWS KMS API requests or validate access against the AWS KMS key policy.

upvoted 1 times

dragosky 5 months, 3 weeks ago

Selected Answer: B

Amazon S3 Bucket Keys reduce the cost of Amazon S3 server-side encryption with AWS Key Management Service (AWS KMS) keys (SSE-KMS). Using a bucket-level key for SSE-KMS can reduce AWS KMS request costs by up to 99 percent by decreasing the request traffic from Amazon S3 to AWS KMS. With a few clicks in the AWS Management Console, and without any changes to your client applications, you can configure your bucket to use an S3 Bucket Key for SSE-KMS encryption on new objects.

upvoted 2 times

djhtoon 6 months, 4 weeks ago

optimize costs associated with encrypting S3 objects without making additional calls to AWS KMS. The answer should be A and it also optimize costs- free.

upvoted 2 times

elmyth 7 months, 2 weeks ago

Selected Answer: A

All S3 are now encrypted by default and free with SSE-S3, but B is also possible.

upvoted 3 times

Sergantus 6 months, 3 weeks ago

Strict data encryption requirements wouldn't be met with SSE-S3

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1010

Topic 1

A company runs multiple workloads on virtual machines (VMs) in an on-premises data center. The company is expanding rapidly. The on-premises data center is not able to scale fast enough to meet business needs. The company wants to migrate the workloads to AWS.

The migration is time sensitive. The company wants to use a lift-and-shift strategy for non-critical workloads.

Which combination of steps will meet these requirements? (Choose three.)

A. Use the AWS Schema Conversion Tool (AWS SCT) to collect data about the VMs.

B. Use AWS Application Migration Service. Install the AWS Replication Agent on the VMs. **Most Voted**

C. Complete the initial replication of the VMs. Launch test instances to perform acceptance tests on the VMs. **Most Voted**

D. Stop all operations on the VMs. Launch a cutover instance. **Most Voted**

E. Use AWS App2Container (A2C) to collect data about the VMs.

F. Use AWS Database Migration Service (AWS DMS) to migrate the VMs.

Correct Answer: BCD

Community vote distribution

BCD (100%)

Comments

LeonSauveterre 4 months, 4 weeks ago

Selected Answer: BCD

Well this is new. Never seen a question like that before.

WHAT YOU SHOULD KNOW:

1. Lift-and-shift means replicating on-premises workloads to AWS with minimal modifications.
 2. A cutover instance is used to direct traffic away from the source region and to the workload in the target region.
- upvoted 4 times

LeonSauveterre 4 months, 4 weeks ago

A - SCT is used for database schema conversions when migrating databases, not for migrating VMs.

B - Application Migration Service is designed for lift-and-shift migrations. It uses the AWS Replication Agent to replicate VMs to AWS.

C - After replication, testing ensures that the workloads function correctly before completing the migration. (If you're wondering why, see this: <https://docs.aws.amazon.com/mgn/latest/ug/launch-cutover-gs.html>)

D - This is a critical step in the migration process. After testing, cutover instances are launched, and the source VMs are stopped to finalize the migration.

E - AWS App2Container is used for containerizing applications, not for migrating VMs.

upvoted 2 times

JA2018 6 months, 1 week ago

Selected Answer: BCD

Why I will choose Options B, C & D:

#B: AWS Application Migration Service (AWS Migration Service)

is the primary tool for a lift-and-shift migration strategy, allowing you to easily replicate on-premises VMs to AWS by installing the Replication Agent on the VMs.

#C: Complete initial replication and launch test instances:

After the initial replication of VMs to AWS, it's crucial to perform acceptance tests on the newly migrated instances to ensure functionality before fully transitioning.

#D: Stop operations and launch a cutover instance:

For a cutover migration, you would stop operations on the on-premises VMs and then launch the migrated instances on AWS to switch over completely

upvoted 2 times

JA2018 6 months, 1 week ago

Why the other options are not suitable:

A and E: AWS Schema Conversion Tool (AWS SCT) and AWS App2Container (A2C):

These tools are primarily used for database migration and containerization respectively, which are not the best fit for a simple lift-and-shift migration of VMs.

F: AWS Database Migration Service (AWS DMS):

While AWS DMS is used for database migration, the question specifies migrating entire VMs, not just databases.

upvoted 1 times

JA2018 6 months, 1 week ago

Key points to remember:

For a lift-and-shift migration, AWS Application Migration Service is the recommended tool.

Always perform testing on migrated instances before fully transitioning to AWS.

A cutover migration involves stopping operations on the source and switching to the new environment.

upvoted 4 times

aragon_saa 6 months, 3 weeks ago

Selected Answer: BCD

Answer is BCD

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1011

Topic 1

A company hosts an application in a private subnet. The company has already integrated the application with Amazon Cognito. The company uses an Amazon Cognito user pool to authenticate users.

The company needs to modify the application so the application can securely store user documents in an Amazon S3 bucket.

Which combination of steps will securely integrate Amazon S3 with the application? (Choose two.)

- A. Create an Amazon Cognito identity pool to generate secure Amazon S3 access tokens for users when they successfully log in. **Most Voted**
- B. Use the existing Amazon Cognito user pool to generate Amazon S3 access tokens for users when they successfully log in.
- C. Create an Amazon S3 VPC endpoint in the same VPC where the company hosts the application. **Most Voted**
- D. Create a NAT gateway in the VPC where the company hosts the application. Assign a policy to the S3 bucket to deny any request that is not initiated from Amazon Cognito.
- E. Attach a policy to the S3 bucket that allows access only from the users' IP addresses.

Correct Answer: AC

Community vote distribution

AC (100%)

Comments

bujuman **Highly Voted** 6 months, 3 weeks ago

Selected Answer: AC

securely integrate Amazon S3 with the application:

<https://docs.aws.amazon.com/cognito/latest/developerguide/identity-pools.html>

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/privatelink-interface-endpoints.html#types-of-vpc-endpoints-for-s3>

upvoted 7 times

FlyingHawk **Most Recent** 4 months, 2 weeks ago

Selected Answer: AC

An Amazon Cognito identity pool allows you to generate temporary AWS credentials (e.g., access tokens) for authenticated users. These credentials can be used to access Amazon S3.

These credentials can be used to securely access AWS services like Amazon S3.

By creating an identity pool, the application can generate secure Amazon S3 access tokens for users after they successfully log in, ensuring that only authenticated users can upload or access documents in the S3 bucket.

upvoted 1 times

FlyingHawk 4 months, 2 weeks ago

C_ A VPC endpoint for Amazon S3 allows private connectivity between the application (hosted in a private subnet) and the S3 bucket. This ensures that data transfer between the application and S3 stays within the AWS network, improving security and performance.

upvoted 1 times

LeonSauveterre 4 months, 4 weeks ago

Selected Answer: AC

A - Amazon Cognito identity pools provide temporary AWS credentials for authenticated users.

B - User pools are for authentication (who the user is). Identity pools are for authorization (what the user can do).

C - Traffic between your VPC and S3 stays within the AWS network, good. This also removes the need for a NAT Gateway for S3 access.

D - Unnecessary and not secure.

E - Users' IP addresses can change frequently (especially mobile users).

upvoted 2 times

viejito 7 months ago

respuesta correcta : A - B

upvoted 2 times

dragossky 5 months, 3 weeks ago

identity pool - authorization

user pool - authentication, in this case something needs to modify docs, so authorization

upvoted 1 times

luther77 6 months, 3 weeks ago

B doesn't make sense here. because user pools are used for authentication, not authorization

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1012

Topic 1

A company has a three-tier web application that processes orders from customers. The web tier consists of Amazon EC2 instances behind an Application Load Balancer. The processing tier consists of EC2 instances. The company decoupled the web tier and processing tier by using Amazon Simple Queue Service (Amazon SQS). The storage layer uses Amazon DynamoDB.

At peak times, some users report order processing delays and hangs. The company has noticed that during these delays, the EC2 instances are running at 100% CPU usage, and the SQS queue fills up. The peak times are variable and unpredictable.


The company needs to improve the performance of the application.

Which solution will meet these requirements?

- A. Use scheduled scaling for Amazon EC2 Auto Scaling to scale out the processing tier instances for the duration of peak usage times. Use the CPU Utilization metric to determine when to scale.
- B. Use Amazon ElastiCache for Redis in front of the DynamoDB backend tier. Use target utilization as a metric to determine when to scale.
- C. Add an Amazon CloudFront distribution to cache the responses for the web tier. Use HTTP latency as a metric to determine when to scale.
- D. Use an Amazon EC2 Auto Scaling target tracking policy to scale out the processing tier instances. Use the ApproximateNumberOfMessages attribute to determine when to scale. **Most Voted**

Correct Answer: D

Community vote distribution

**Comments**

LeonSauveterre 4 months, 4 weeks ago

Selected Answer: D

Just scale base on how many tasks left in the queue. CPU utilization is not an accurate metric. There could be just one task but it takes quite a while to process, in which case scaling is unnecessary.

upvoted 2 times

FlvinnHawk 5 months, 1 week ago

tyingham 5 months, 1 week ago

Selected Answer: D

The bottleneck is at the processing tier, B and C are incorrect, processing tier instances are running at 100% CPU usage, and the SQS queue fills up. Because the peak times are variable and unpredictable, we should use the EC2 Auto Scaling target tracking policy to scale out the processing tier instances based on the size of the queue, A uses scheduled scaling for Amazon EC2 Auto Scaling is incorrect.

upvoted 2 times

GOTJ 5 months, 1 week ago

Selected Answer: D

Check this out:

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-using-sqs-queue.html>

upvoted 1 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: D

The bottleneck is from the processing tier.

upvoted 3 times

dragosky 5 months, 3 weeks ago

Selected Answer: B

Yes, ElastiCache can be used with DynamoDB to improve the performance and scalability of read-heavy or frequently accessed workloads. Amazon ElastiCache is a fully managed in-memory data store that supports popular caching engines such as Memcached and Redis.

upvoted 1 times

Salilgen 4 months, 3 weeks ago

The problem is from application tier

upvoted 1 times

GOTJ 5 months, 1 week ago

Agree. But this "SQS queue fill up" (remember, SQS is meant for decouple web and processing tiers) suggest that the bottleneck is produced by the processing CPU bottleneck because web servers are sending more requests to the queue than the processing servers are able to process from the queue. As the first approach, I would try option D

upvoted 1 times

EllenLiu 5 months, 2 weeks ago

DynamoDB Accelerator is cache for dynamoDB

upvoted 1 times

aragon_saa 7 months, 3 weeks ago

Selected Answer: D

Answer is D

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1013

Topic 1

A company's production environment consists of Amazon EC2 On-Demand Instances that run constantly between Monday and Saturday. The instances must run for only 12 hours on Sunday and cannot tolerate interruptions. The company wants to cost-optimize the production environment.

Which solution will meet these requirements MOST cost-effectively?

- A. Purchase Scheduled Reserved Instances for the EC2 instances that run for only 12 hours on Sunday. Purchase Standard Reserved Instances for the EC2 instances that run constantly between Monday and Saturday. **Most Voted**
- B. Purchase Convertible Reserved Instances for the EC2 instances that run for only 12 hours on Sunday. Purchase Standard Reserved Instances for the EC2 instances that run constantly between Monday and Saturday.
- C. Use Spot Instances for the EC2 instances that run for only 12 hours on Sunday. Purchase Standard Reserved Instances for the EC2 instances that run constantly between Monday and Saturday.
- D. Use Spot Instances for the EC2 instances that run for only 12 hours on Sunday. Purchase Convertible Reserved Instances for the EC2 instances that run constantly between Monday and Saturday.

Correct Answer: A

Community vote distribution

A (100%)

Comments

LeonSauveterre 4 months, 4 weeks ago

Selected Answer: A

Scheduled Reserved Instance: Perfect for workloads with predictable and recurring schedules (like 12 hours on Sunday). This is cost-efficient and ensures uninterrupted instances.

Convertible Reserved Instance: Allows flexibility to change instance types, sizes, or OS during the term. However, they are not ideal for short, predictable schedules like 12 hours on Sunday. Scheduled Reserved Instances are a better fit.

Spot Instances can be interrupted at any time. They do not guarantee uninterrupted operations.

upvoted 3 times

FlyingHawk 5 months, 1 week ago

Selected Answer: A

as the instance cannot tolerate interruptions, so spot instances are out, C and D are incorrect, leave only A and B, Scheduled RI or Convertible RI, based on the AWS doc, Purchasing Reserved Instances on a recurring schedule lets you pay for compute capacity by the hour and obtain a capacity reservation ahead of time for only the time periods you will need it. as it only 12 hours on Sunday, so A will be cheaper than B
<https://aws.amazon.com/ec2/pricing/reserved-instances/pricing/>.
upvoted 1 times

FlyingHawk 5 months, 1 week ago

Not the minimum commit for RI is one year, if you want the flexibility of change instance type within one year, B might be a better choice, however, considering the flexibility, saving plan for computing is a recommended solution over RIs, I guess this question should not appear in the exam if AWS wants to advocate the saving plan.
upvoted 1 times

GOTJ 5 months, 1 week ago

Selected Answer: A

"A" is the most straightforward option. However, I'm afraid this feature could be eventually removed and substituted by On-Demand Capacity Reservations and Savings Plans. If you try to use scheduled reservation in the EC2 console, you'll get the following banner:

"We recommend Savings Plans over Reserved Instances. Savings Plans are the easiest and most flexible way to save money on your AWS compute costs and offer lower prices (up to 72% off) just like Reserved Instances. To learn more and get started with Savings Plans click here"
upvoted 1 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: A

Scheduled reserved instances are purchase of reserved Instances on a recurring schedule
<https://aws.amazon.com/ec2/pricing/reserved-instances/buyer/>
upvoted 1 times

bujuman 6 months, 3 weeks ago

Selected Answer: A

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-reserved-instances.html>
upvoted 3 times

8621a7c 7 months, 3 weeks ago

Selected Answer: A

standard reserved instance is cheaper than convertible reserved instance
upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1014

Topic 1

A digital image processing company wants to migrate its on-premises monolithic application to the AWS Cloud. The company processes thousands of images and generates large files as part of the processing workflow.

The company needs a solution to manage the growing number of image processing jobs. The solution must also reduce the manual tasks in the image processing workflow. The company does not want to manage the underlying infrastructure of the solution.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 Spot Instances to process the images. Configure Amazon Simple Queue Service (Amazon SQS) to orchestrate the workflow. Store the processed files in Amazon Elastic File System (Amazon EFS).
- B. Use AWS Batch jobs to process the images. Use AWS Step Functions to orchestrate the workflow. Store the processed files in an Amazon S3 bucket. **Most Voted**
- C. Use AWS Lambda functions and Amazon EC2 Spot Instances to process the images. Store the processed files in Amazon FSx.
- D. Deploy a group of Amazon EC2 instances to process the images. Use AWS Step Functions to orchestrate the workflow. Store the processed files in an Amazon Elastic Block Store (Amazon EBS) volume.

Correct Answer: B

Community vote distribution

B (100%)

Comments

MortisG 2 weeks, 4 days ago

Selected Answer: B

All other options require managing infrastructure. AWS Batch handles the processing and Step Functions the workflow. Best solutions for what's solicited.

upvoted 1 times

ElI89 2 months ago

Selected Answer: B

B- because all other answers mention EC2 which goes against the company not wanting to manage the underlying infrastructure.
upvoted 2 times

LeonSauveterre 4 months, 4 weeks ago

Selected Answer: B

A & C - Spot Instances are cost-effective but are subject to interruptions. "The solution must also reduce the manual tasks."
B - AWS Batch is a fully managed service for batch processing, ideal for workloads like image processing. It scales automatically, handles job scheduling, and eliminates the need to manage infrastructure. Step Functions provides robust orchestration with built-in support for retries, error handling, and dependencies, reducing manual tasks in the workflow.
D - Using EC2 instances requires managing the underlying infrastructure. Also, EBS is not as scalable or cost-effective as S3 when you want to store large, growing numbers of files.

upvoted 3 times

trinh_le 6 months, 1 week ago

Selected Answer: B

A. Use Amazon Elastic Container Service (Amazon ECS) complexity
C. Amazon FSx is not an ideal compared to s3
D. Use a group of Amazon EC2 instances do not scalable

upvoted 4 times

aragon_saa 7 months, 3 weeks ago

Selected Answer: B

Answer is B

upvoted 4 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1015

Topic 1

A company's image-hosting website gives users around the world the ability to upload, view, and download images from their mobile devices. The company currently hosts the static website in an Amazon S3 bucket.

Because of the website's growing popularity, the website's performance has decreased. Users have reported latency issues when they upload and download images.

The company must improve the performance of the website.

Which solution will meet these requirements with the LEAST implementation effort?

- A. Configure an Amazon CloudFront distribution for the S3 bucket to improve the download performance. Enable S3 Transfer Acceleration to improve the upload performance. **Most Voted**
- B. Configure Amazon EC2 instances of the right sizes in multiple AWS Regions. Migrate the application to the EC2 instances. Use an Application Load Balancer to distribute the website traffic equally among the EC2 instances. Configure AWS Global Accelerator to address global demand with low latency.
- C. Configure an Amazon CloudFront distribution that uses the S3 bucket as an origin to improve the download performance. Configure the application to use CloudFront to upload images to improve the upload performance. Create S3 buckets in multiple AWS Regions. Configure replication rules for the buckets to replicate users' data based on the users' location. Redirect downloads to the S3 bucket that is closest to each user's location.
- D. Configure AWS Global Accelerator for the S3 bucket to improve network performance. Create an endpoint for the application to use Global Accelerator instead of the S3 bucket.

Correct Answer: A*Community vote distribution*

A (100%)

Comments**FlyingHawk** 4 months, 2 weeks ago**Selected Answer: A**

A - By configuring a CloudFront distribution for the S3 bucket, the company can significantly improve download performance for users by serving images from the nearest edge location. S3 Transfer Acceleration uses CloudFront's globally distributed edge

access by serving images from the nearest edge location. CDN transfer acceleration uses Amazon's globally distributed edge locations to accelerate upload performance. It optimizes the path between the user's device and the S3 bucket, reducing latency and improving upload speeds.

upvoted 2 times

FlyingHawk 4 months, 2 weeks ago

Option C (Configure CloudFront for downloads, use CloudFront for uploads, create S3 buckets in multiple Regions, and configure replication rules):

While this option improves performance, it requires significant effort to configure multiple S3 buckets, replication rules, and redirection logic. It is more complex than necessary for the company's requirements.

upvoted 2 times

Denise123 5 months, 1 week ago

Selected Answer: A

Obvious.

upvoted 1 times

aragon_saa 7 months, 3 weeks ago

Selected Answer: A

Answer is A

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1016

Topic 1

A company runs an application in a private subnet behind an Application Load Balancer (ALB) in a VPC. The VPC has a NAT gateway and an internet gateway. The application calls the Amazon S3 API to store objects.

According to the company's security policy, traffic from the application must not travel across the internet.

Which solution will meet these requirements MOST cost-effectively?

- A. Configure an S3 interface endpoint. Create a security group that allows outbound traffic to Amazon S3.
- B. Configure an S3 gateway endpoint. Update the VPC route table to use the endpoint. **Most Voted**
- C. Configure an S3 bucket policy to allow traffic from the Elastic IP address that is assigned to the NAT gateway.
- D. Create a second NAT gateway in the same subnet where the legacy application is deployed. Update the VPC route table to use the second NAT gateway.

Correct Answer: B*Community vote distribution*

B (100%)

Comments**sOI852POL** **Highly Voted** 7 months, 2 weeks ago**Selected Answer: B**

S3 gateway endpoint

<https://docs.aws.amazon.com/vpc/latest/privatelink/vpc-endpoints-s3.html>

upvoted 5 times

GOTJ **Most Recent** 5 months, 1 week ago**Selected Answer: B**

Here's another document that support option "B" as the most cost-effectively:

<https://aws.amazon.com/blogs/architecture/choosing-your-vpc-endpoint-strategy-for-amazon-s3/>

upvoted 2 times

trinh_le 6 months, 1 week ago

Selected Answer: B

Keyword: ec2 connect s3 -> vpc endpoint or s3 endpoint
upvoted 3 times

aragon_saa 7 months, 3 weeks ago

Selected Answer: B

Answer is B
upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1017


Topic 1

A company has an application that runs on an Amazon Elastic Kubernetes Service (Amazon EKS) cluster on Amazon EC2 instances. The application has a UI that uses Amazon DynamoDB and data services that use Amazon S3 as part of the application deployment.

The company must ensure that the EKS Pods for the UI can access only Amazon DynamoDB and that the EKS Pods for the data services can access only Amazon S3. The company uses AWS Identity and Access Management (IAM).

Which solution meets these requirements?

- A. Create separate IAM policies for Amazon S3 and DynamoDB access with the required permissions. Attach both IAM policies to the EC2 instance profile. Use role-based access control (RBAC) to control access to Amazon S3 or DynamoDB for the respective EKS Pods.
- B. Create separate IAM policies for Amazon S3 and DynamoDB access with the required permissions. Attach the Amazon S3 IAM policy directly to the EKS Pods for the data services and the DynamoDB policy to the EKS Pods for the UI.
- C. Create separate Kubernetes service accounts for the UI and data services to assume an IAM role. Attach the AmazonS3FullAccess policy to the data services account and the AmazonDynamoDBFullAccess policy to the UI service account. **Most Voted**
- D. Create separate Kubernetes service accounts for the UI and data services to assume an IAM role. Use IAM Role for Service Accounts (IRSA) to provide access to the EKS Pods for the UI to Amazon S3 and the EKS Pods for the data services to DynamoDB.

Correct Answer: C*Community vote distribution***Comments**

jingen11 **Highly Voted** 7 months, 3 weeks ago

Selected Answer: D

<https://docs.aws.amazon.com/eks/latest/userguide/service-accounts.html#service-accounts-iam>
upvoted 8 times

LeonSauveterre Highly Voted 4 months, 4 weeks ago

Selected Answer: C

A - This is a major security flaw. All Pods running on the EC2 instances would inherit both the S3 and DynamoDB permissions.
B - Directly attaching IAM policies to Pods is not a valid AWS mechanism. You should use roles.
C - Not recommended unless we throw least privilege principles to the wind. BUT this is the only feasible option here, so.
D - This option mixes up the permissions. It gives UI Pods access to S3 and data service Pods access to DynamoDB, which is the OPPOSITE of the requirement.

upvoted 5 times

FlyingHawk Most Recent 4 months, 2 weeks ago

Selected Answer: C

IRSA allows Kubernetes service accounts to assume IAM roles, enabling fine-grained access control for EKS Pods. This ensures that each Pod can access only the AWS resources it needs.

Attach the AmazonDynamoDBFullAccess policy to the service account for the UI Pods, allowing them to access only DynamoDB. Attach the AmazonS3FullAccess policy to the service account for the data services Pods, allowing them to access only S3.

upvoted 3 times

BugsyWarribwoy 5 months ago

Selected Answer: C

D is sneakily WRONG because it swaps the services incorrectly. Pay attention. UWC.

upvoted 5 times

GOTJ 5 months, 1 week ago

Selected Answer: C

Even though there are a couple of comments rightfully discarding "D" because the UI --> S3/DataServices --> DynamoDB swap, I found an AWS document claiming that Kubernetes Service accounts and IAM Role for Service Accounts combination should be the right answer:

<https://docs.aws.amazon.com/eks/latest/userguide/iam-roles-for-service-accounts.html>

Having this in mind, I'm also discarding "D" (right reasoning, wrong scenario), as well "A" and "B". I didn't like the "full access" policy, but is technically correct, so my vote goes to "C"

upvoted 3 times

trinh_le 6 months, 1 week ago

Selected Answer: C

* A. Attach both IAM policies to the EC2 instance profile. Does not separate each EKS Pod

* B. Aws does not support Attach the Amazon S3 IAM policy directly to the EKS Pods

*

* D. provide access "UI to Amazon S3" and "data services to DynamoDB" => it does not meet requirements

upvoted 3 times

rosanna 6 months, 1 week ago

Selected Answer: C

The answer is C as they're switching data pods with DynamoDB service and vice versa (configure the wrong resources)

upvoted 2 times

tm1000000 7 months, 2 weeks ago

answer is D

upvoted 1 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1018

Topic 1

A company needs to give a globally distributed development team secure access to the company's AWS resources in a way that complies with security policies.

The company currently uses an on-premises Active Directory for internal authentication. The company uses AWS Organizations to manage multiple AWS accounts that support multiple projects.


The company needs a solution to integrate with the existing infrastructure to provide centralized identity management and access control.

Which solution will meet these requirements with the LEAST operational overhead?

- A. Set up AWS Directory Service to create an AWS managed Microsoft Active Directory on AWS. Establish a trust relationship with the on-premises Active Directory. Use IAM roles that are assigned to Active Directory groups to access AWS resources within the company's AWS accounts.
- B. Create an IAM user for each developer. Manually manage permissions for each IAM user based on each user's involvement with each project. Enforce multi-factor authentication (MFA) as an additional layer of security.
- C. Use AD Connector in AWS Directory Service to connect to the on-premises Active Directory. Integrate AD Connector with AWS IAM Identity Center. Configure permissions sets to give each AD group access to specific AWS accounts and resources. **Most Voted**
- D. Use Amazon Cognito to deploy an identity federation solution. Integrate the identity federation solution with the on-premises Active Directory. Use Amazon Cognito to provide access tokens for developers to access AWS accounts and resources.

Correct Answer: C

Community vote distribution



Comments

LeonSauveterre 4 months, 4 weeks ago

Selected Answer: C

A - On-premises Active Directory already exists.

B - "Manually" ? Oh hell no. But MFA is something I guess.

C - AD Connector seamlessly connects AWS to the on-premises Active Directory without the need to synchronize or replicate the directory. Identity Center (formerly AWS SSO) allows centralized access management across AWS accounts in an AWS Organizations setup. Permissions sets can be configured to map Active Directory groups to specific AWS accounts and resources, making access control easy and secure.

D - Amazon Cognito is better suited for application-level identity management (like customer-facing apps), not for internal teams working across multiple AWS accounts.

upvoted 3 times

GOTJ 5 months, 2 weeks ago

Selected Answer: A

Why not "A"? Check out the note of this link: https://docs.aws.amazon.com/directoryservice/latest/admin-guide/directory_ad_connector.html

AD Connector cannot be shared with other AWS accounts. If this is a requirement, consider using AWS Managed Microsoft AD to Share your AWS Managed Microsoft AD. AD Connector is also not multi-VPC aware, which means that AWS applications like WorkSpaces are required to be provisioned into the same VPC as your AD Connector.

And I think managing multiple aws accounts is, indeed, a requirement

upvoted 3 times

FlyingHawk 4 months, 2 weeks ago

While this option provides centralized identity management, it requires setting up a separate AWS managed Microsoft AD, which increases operational overhead. A trust relationship is also more complex to configure than AD Connector.

upvoted 1 times

GOTJ 4 months ago

I guess so, but... is the alternative configure one AD connector per account more operational friendly?

upvoted 1 times

trinh_le 6 months, 1 week ago

Selected Answer: C

C: AD Connector allows AWS to use the on-premises Active Directory for authentication without replicating directory data to AWS.
• Maintains centralized identity management in the on-premises directory, adhering to the company's security policies.

A: Higher Overhead: Requires creating and maintaining a separate managed Active Directory instance in AWS.

B: Scalability Issues: Manually creating and managing IAM users for a globally distributed team is cumbersome.

D: Unnecessary Federation Layer: Cognito is more suited for customer identity use cases rather than managing internal developer access to AWS resources.

upvoted 2 times

78b9037 6 months, 2 weeks ago

Selected Answer: C

Once set up, it requires minimal ongoing management. User provisioning and deprovisioning are handled through the existing Active Directory processes.

upvoted 2 times

xekiva3329 7 months, 3 weeks ago

Selected Answer: C

answer C

upvoted 2 times

aragon_saa 8 months ago

Selected Answer: C

Answer is C

upvoted 2 times

Exam AWS Certified Solutions Architect - Associate SAA-C03 All Actual Questions

Question #1019

Topic 1


A company is developing an application in the AWS Cloud. The application's HTTP API contains critical information that is published in Amazon API Gateway. The critical information must be accessible from only a limited set of trusted IP addresses that belong to the company's internal network.

Which solution will meet these requirements?

- A. Set up an API Gateway private integration to restrict access to a predefined set of IP addresses.
- B. Create a resource policy for the API that denies access to any IP address that is not specifically allowed. **Most Voted**
- C. Directly deploy the API in a private subnet. Create a network ACL. Set up rules to allow the traffic from specific IP addresses.
- D. Modify the security group that is attached to API Gateway to allow inbound traffic from only the trusted IP addresses.

Correct Answer: B

Community vote distribution



Comments

Jassn3 **Highly Voted** 8 months ago

OMG finally I have reached here
upvoted 26 times

Bwhizzy **Highly Voted** 7 months, 3 weeks ago

Finally made it. Congratulations to everyone who got here. I know the journey has been hard and long. best of luck
upvoted 15 times

whoevd2eww3322d **Most Recent** 4 months ago

□ I cleared the AWS Solutions Architect – Associate exam!

For those on the same journey, many exam questions are from or related to these 1019 questions. Go through them twice for a solid prep.
Also, check out Shaping Pixel's YouTube channel—focus only on the latest videos as most questions overlap.

You've got this—good luck! □

upvoted 5 times

techghost 4 months, 2 weeks ago

Selected Answer: B

Got to the end, worth it and great contributions

upvoted 1 times

LeonSauveterre 4 months, 4 weeks ago

Selected Answer: B

A - A private integration is for connecting API Gateway to private resources within VPCs, such as EC2 instances or Lambda functions. It doesn't restrict access based on IP addresses. It controls how the API interacts with private resources.

B - This is the most direct way to restrict access at the API Gateway level.

C - Deploying the API in a private subnet would require users to connect to your VPC (e.g., via VPN or Direct Connect). This is not what the question implies (the API should be publicly accessible, but restricted by IP).

D - API Gateway doesn't even use security groups because it's a regional managed service, not directly tied to VPCs like EC2 instances.

upvoted 4 times

LeonSauveterre 4 months, 4 weeks ago

Now people let's go back to square one and nail this frickin' exam!

upvoted 4 times

jfedotov 5 months ago

Selected Answer: B

The correct answer should be B, but the question is misleading. This is because HTTP APIs do not support resource policies; only REST APIs do.

upvoted 1 times

Denise123 5 months, 1 week ago

Selected Answer: B

I can't believe that I made it! Congrats everyone! Time to restart from the Q1 to polish, and good luck for the exam!

upvoted 1 times

EllenLiu 5 months, 2 weeks ago

Selected Answer: B

<https://docs.aws.amazon.com/apigateway/latest/developerguide/apigateway-resource-policies-examples.html#apigateway-resource-policies-source-ip-address-example>

upvoted 3 times

JA2018 6 months, 1 week ago

Selected Answer: D

yes, finally, the finishing point.....

Why not D?

By modifying the security group attached to API Gateway, you can explicitly define which IP addresses are allowed to access the API, effectively restricting access to only the trusted internal network IPs

upvoted 2 times

trinh_le 6 months, 1 week ago

Selected Answer: B

B: api gateway allows configure resource policy to restricted IP

A: private integration-> private api

C: private subnet -> private api

D: security group does not support api gateway

upvoted 5 times

Sergantus 6 months, 3 weeks ago

Congrats to everyone who made it here. On to the next step, congrats

Congrats to everyone who made it here. On to the next step, gents.

upvoted 6 times

Ben_88 7 months, 1 week ago

And now it's time to start all over again to see if you can do better this time . good luck lads !

upvoted 5 times

xekiva3329 7 months, 3 weeks ago

Selected Answer: B

answer B

upvoted 5 times